RESEARCH-SCHEDULING DIRECTIVE — FEATURE (§11.4.213). The remainder of this prompt DESCRIBES a feature to research and plan. It MUST be turned into a real, fully-populated, fully-synced workable item — never merely acknowledged, never silently dropped (§11.4.202 / §11.4.197). Unlike BUG / TASK / ISSUE, FEATURE SCHEDULES rather than synchronously executes: create the item NOW through the project's feature-scheduling engine (`constitution/scripts/feature/schedule_feature_research.sh --description <text>`), which (1) creates a Type=Task + Status=Queued workable item with a stable auto-incremented ticket id in the workable-items SQLite single-source-of-truth (§11.4.93 / §11.4.95 / §11.4.54) — by INVOKING the SAME §11.4.202 `report_item.sh` engine, never a duplicate implementation — whose §11.4.148 / §11.4.171 comprehensive description embeds the full research-and-planning WORK PROGRAM verbatim: (a) deep web research (articles / guides / scientific papers / open-source codebases / examples) on the best way to incorporate this feature (§11.4.8 / §11.4.99 / §11.4.150); (b) systematic-debug (§11.4.102) of ALL data obtained to enumerate EVERY weak spot / gap / danger-zone / inconsistency / imperfection and design a risk-free rock-solid solution for EACH one; (c) ALWAYS enterprise-grade, bleeding-edge, innovative — never less; (d) MANDATORY exhaustive documentation — plans down to lines-of-code + micro-POCs, diagrams, illustrations, SQL definitions, templates, every relevant material (§11.4.65 / §11.4.73); (e) investigate + plan REUSE of existing vasic-digital / HelixDevelopment submodules/components BEFORE proposing a rewrite, extending them freely while keeping them decoupled / reusable (§11.4.28 / §11.4.74 / §11.4.177); (f) plan from the TESTING point of view — every supported test type + Challenges + full HelixQA bank coverage (§11.4.27 / §11.4.169); (g) divide into fine-grained, nano-detailed phases / tasks / subtasks ready to drive integration / implementation / wiring / testing / scaling; (h) plan for enterprise scalability + maximal performance; (i) when there is a UI/UX surface, produce full wireframes / diagrams / design files authored via OpenDesign (§11.4.162 / §11.4.190); (j) plan full CodeGraph integration with regular/real-time index sync (§11.4.78 / §11.4.79 / §11.4.80); (k) create fully-detailed follow-on workable items synced in REAL TIME to the SQLite SSoT + every derived document + every connected external tracker (§11.4.148 D5), honestly SKIPPING any absent tracker

(§11.4.10) — NEVER a faked push. (2) enqueues the item in the durable `docs/requests/feature_queue.md` queue so it can NEVER be silently dropped or forgotten (mirrors the BACKGROUND durable queue, §11.4.140); (3) regenerates every derived document from the DB (§11.4.12 / §11.4.53 / §11.4.65 / §11.4.106) and pushes to every configured external tracker exactly as §11.4.202 does. Report back the assigned ticket id, the research-doc destination path, the sync verdict, and each tracker's verdict with its captured-evidence path (§11.4.5 / §11.4.69). THEN — the scheduled research/planning/documentation work itself is EXECUTED LATER by the standing autonomous loop (§11.4.87 / §11.4.94 / §11.4.97 / §11.4.103 / §11.4.126) when it claims the item, driven to a genuinely COMPLETED-and-wired or explicitly evidence-backed CLOSED terminal state under §11.4.197 — it is NEVER performed synchronously inline by this directive, and NEVER left un-wired in the backlog.

$ARGUMENTS