

terse_output_guidance.md – § . . Measure M (output-token reduction)

Field	Value
Revision	1
Last modified	2026-06-09T00:00:00Z
Status	active
Authority	constitution submodule, universal (§11.4.17). Measure M5 of §11.4.141 — output is ~5× input price, so terseness is cost-leveraged.

Reusable instruction snippet. Paste verbatim into a subagent dispatch / system instruction / agent-config to apply the §11.4.141 low-effort, concise-output discipline. **Presentation-only** — captured-evidence artefacts (§11.4.5/ §11.4.69) are UNCHANGED; only the conductor's narration shrinks. `effort:` "low" is confined to the **mechanical** allowlist (§11.4.141 M2 / `subagent_model_tiering.md`) so **no reasoning / verdict path is degraded**.

The snippet (paste verbatim)

OUTPUT DISCIPLINE (§11.4.141 M5 – terse, cost-leveraged; presentation-only):

- Be concise. One line per milestone; no restatement of unchanged governance or rules.
- Do NOT echo file contents you only read; cite path:line instead.
- Final report: only what the caller needs to act – paths (absolute), the load-bearing facts, and the verdict. No recap of code you merely read.
- Structured output (JSON/TSV) wherever a sink consumes it – never prose a machine will parse.
- effort:"low" applies ONLY to mechanical work (search / grep / status / doc-export / parse / file-presence – the §11.4.141 M2 mechanical allowlist). NEVER lower effort on a reasoning / fix-design / code-review / verdict path.
- Terseness changes the NARRATION, never the EVIDENCE: every PASS still cites its captured artefact (§11.4.5 / §11.4.69); ab_pass_with_evidence is unchanged; no analysis is dropped.
- No emojis. No filler ("Sure!", "Great question", "Let me..."). State the result.

Why it cannot lower quality

Concern	Why M5 is safe
"Less output → less rigour?"	Output is <i>narration</i> . Evidence is captured to disk regardless (§11.4.5/§11.4.69).
"effort:low → weaker reasoning?"	effort:low is mechanical-only (M2 allowlist). Reasoning/verdict paths keep full effort.
"Could it drop a PASS's evidence?"	No — ab_pass_with_evidence and the captured-artefact requirement are untouched; a terse PASS still cites its artefact path.

A config that suppresses *evidence capture* (not just narration) is the violation the §1.1 mutation exploits (the feature test's evidence assertion FAILs).

1 1 4 2 8

Decoupling / reuse (\$. .)

The snippet names no project, package, model id, or path. Any consuming project pastes it into its own dispatch flow as-is.

Cross-references

- §11.4.141 (M5) · §11.4.141 M2 / `subagent_model_tiering.md` (the mechanical allowlist that bounds `effort: low`) · §11.4.5 / §11.4.69 (evidence unchanged) · §1.1 (suppress-evidence mutation) · §11.4.6 (no-guessing — state results, don't hedge).