

AGENTS

Helix Constitution — Universal AGENTS.md

Field	Value
Revision	33
Created	2026-05-14
Last modified	2026-06-21T00:00:00Z
Status	active
Status summary	Added §11.4.163 mirror — Universal Media Validation & Verification Mandate: every recorded artifact MUST pass MEDIA VALIDATION pipeline before acceptance (OCR/transcription/text parsing vs SPECIFY-phase patterns, self-validated analyzer, structured verdict with pinpoint on FAIL, post-recording+real-time triggers, paired §1.1 mutation). Added §11.4.164 mirror — Universal Constitution Auto-Propagation & Hook System: post_update_hook.sh detects/registers/installs changed skills/MCP/hooks/scripts; consumer invokes after every constitutional pull. Added §11.4.165 mirror — Universal Independent Verification Agent Mandate: every output passes INDEPENDENT verifier (§11.4.70/.20), iterates to zero-finding GO (§11.4.134). All Classification: universal (§11.4.17). Prior round: Added §11.4.162 mirror — OpenDesign UI design system mandate (User mandate 2026-06-21): every project producing user-facing interfaces MUST use OpenDesign as the mandatory UI design-and-refinement system — NOT ad hoc CSS or one-off design tools; install as a project dependency and use its design tokens/themes for color palette (light+dark), typography, spacing, and component-level tokens; extend upstream per §11.4.74 for missing patterns; every UI component ships light+dark variants with project brand colors from canonical assets; elements MUST NOT overlap or overlay labels; all UI changes covered by standard test types including visual regression. Classification: universal (§11.4.17). Prior round: Added §11.4.161 mirror — Rootless container runtime mandate (User mandate 2026-06-21): every project MUST use Podman in rootless mode (or equivalent rootless container runtime) for ALL containerized workloads — Docker in rootful mode, sudo, or any escalation to root is FORBIDDEN unless the target platform has no rootless option AND that constraint is documented per §11.4.112; the vasic-digital/containers submodule (§11.4.76) MUST be used as the sole container orchestration layer — no ad-hoc docker/podman commands outside pkg/boot/pkg/compose/pkg/health; missing capabilities extend to container Submodule upstream per §11.4.74; integration tests boot infra on-demand via the Submodule. Classification: universal (§11.4.17). Prior round: Added §11.4.160 mirror — Vision-verified recording + HelixQA bridge mandate (User mandate 2026-06-21): every video recording for feature/QA evidence MUST be processed through a vision/OCR pipeline that reads on-screen content and confirms expected results BEFORE

Field	Value
	acceptance; the recording system MUST provide a bridge feeding captured frames to HelixQA's test infrastructure (or equivalent) for content verification — automated read-the-screen against specified expected patterns; the bridge MUST capture frames at ≤ 5s intervals, run OCR/vision analysis with self-validated analyzer (§11.4.107(10)), compare against SPECIFY-phase patterns, produce per-frame PASS/FAIL with evidence path, surface failures immediately for re-record per §11.4.159(L). Honest boundary: vision verification confirms on-screen content — does NOT replace §11.4.5 quality analysis nor §11.4.108 runtime-signature; FLAG_SECURE surfaces use the §11.4.117 proxy oracle. Classification: universal (§11.4.17). Prior round: Added §11.4.153 mirror — Crashlytics recorded-data continuous monitoring + systematic-debug + regression-test-coverage mandate (User mandate 2026-06-13): every project with Firebase Crashlytics enabled/wired MUST continuously monitor ALL four recorded surfaces (fatal crashes, ANRs, performance traces, AND non-fatals), systematic-debug each (reproduce-before-fix §11.4.102/.115), fix/improve, and cover every closure with a permanent §11.4.135 regression guard (RED-on-broken→GREEN-on-fixed) + a closure log citing the console issue id/URL + the validation/verification test paths; regular cadence (the §11.4.47 five-trigger set), no false results, no bluff — a console-resolved Crashlytics issue with no falsifiable regression test is FORBIDDEN (the silent-recurrence vector). STRENGTHENS+COMPLETES §11.4.47 (review/dedup finds it; §11.4.152 fixes it + proves it stays fixed). Project-level reference instantiation = the Lava client's §6.O + §6.AC. Composes §11.4/.1/.5/.6/.9/.34/.40/.43/.47/.69/.90/.102/.107/.108/.115/.118/.123/.130/.135/.138/.15 §1.1. Propagation gate CM-COVENANT-114-152-PROPAGATION + recommended gate CM-CRASHLYTICS-ISSUE-FULLY-COVERED. Classification: universal (§11.4.17). Prior round: Added §11.4.151 mirror — Project-prefixed release-tag/version-naming mandate (User mandate 2026-06-12): every release tag AND version name on the main repo AND every owned submodule MUST be prefixed <PREFIX>-<version> (e.g. helix_translate-1.0.0-dev-0.0.1); prefix resolution order (1) HELIX_RELEASE_PREFIX from .env (authoritative, git-ignored §11.4.30, documented in tracked .env.example §11.4.77) else (2) lowercased snake_case project-root dir name §11.4.29; SAME prefix across main + all owned submodules in one release = greppable across every repo; version codes increment monotonically. Composes §2.1/ §11.4.28/.29/.30/.40/.77/.113/.126. Propagation gate CM-COVENANT-114-151-PROPAGATION + recommended gate CM-RELEASE-PREFIX-NAMING. Classification: universal (§11.4.17). Prior round: Added §11.4.148 + §11.4.149 mirrors — workable-item integrity + testing-diary family (User mandate 2026-06-10). §11.4.148 BINDS+STRENGTHENS §11.4.15/.16/.21/.54/.91/.93/.95/.106 into one DB[?]docs[?]external-tracker integrity contract: (D1) no item without a valid status+type+stable id on ALL three surfaces; (D2) comprehensive structured description (what/manifest/repro/acceptance); (D3) BLOCKED items carry WHY + enumerated unblock CHOICES (tightens §11.4.21; Blocked/BLOCKED = documented alias of canonical Operator-blocked); (D4) regular never-missed bidirectional DB[?]docs[?]tracker sync, §11.4.86-fingerprinted + §11.4.106 docs_chain-bound; (D5) generic idempotent external-tracker push (statuses/types/assignee-from-env/sub-tasks, sink-side proof §11.4.69, machinery project-agnostic §11.4.28). §11.4.149 per-workable item TESTING DIARY (date/time + tested-by + result + observations + action+why), distinct from item_history/Reopens: append-only test_diary table with PASS-required-evidence schema constraint + in-depth diary + derived summary VIEW + four-format per-item exports §11.4.86-fingerprinted §11.4.106-bound + external-tracker SUB-TASK {TODO In-progress Completed} lifecycle + minimal-LLM deterministic tooling 100% test-covered §11.4.27. Both Classification: universal (§11.4.17); project-specific dual per-display brightness/auto-dim/font-size landed ONLY in the consumer layer (RK3588/Presenter), NOT here. Propagation gates CM-COVENANT-114-148-

Field	Value
	PROPAGATION + CM-COVENANT-114-149-PROPAGATION + paired §1.1 mutations. Prior round: Added §11.4.147 mirror — Crashed-agent respawn-until-complete + no-work-lo registry mandate (User mandate 2026-06-10): every dispatched agent/subagent MUST b tracked through its full lifecycle so a crash NEVER loses/forgets/corrupts its work (cras ≠ done, §11.4.6) — (a) durable append-only agent REGISTRY (status CLOSED SET {dispatched in-flight crashed respawned complete}) on the §11.4.116 sync substrate, SINGLE SOURCE OF TRUTH for “is this work owed?”; (b) mechanical CRASH-DETECTION (transient rate-limit/socket-close) → flip to crashed + keep OPEN → RESPAWN until complete autonomously per §11.4.101 with ALREADY-DEFINED backoff; (c) PARTIAL-STATE preserve → §11.4.84-check → resume-or-clean-restart (§9.2 backup) in a per-agent worktree; (d) the §11.4.87/.94/.97/.126 loop done-condition MUST NOT read satisfied while any entry is non-complete. Forensic FACT this session: 5 subagents killed by rate-limit + 2 by socket-close + one PARTIAL dual-sliders edit, re-dispatched by pure conductor vigilance with NO mechanical guarantee; agent-side analogue of §11.4.144. Composes §11.4.6/.58/.84/.87/.94/.97/.101/.116/.102/.108/.125/.142/.40/.128/.144/§9.2/§1.1. Classification: universal (§11.4.17). Prior round: Added §11.4.144 mirror — Tracked/recorded-device availability-following mandate (User mandate 2026-06-10): every device the project is tracking / following / recording (test / debug / manual-testing, across every reachable transport — USB / wireless ADB / SSH / serial / network introspection API) MUST be availability-FOLLOWED — connection state continuously monitored, any drop handled, never silently abandoned and never presented as a continuous recording; the §11.4.128 always-on recorder KNOWS a tracked device is absent (its loo guards on the reachable state) but, lacking a following discipline, merely spins idle (no data, no offline event, no resume, no escalation) → silent corpus hole = §11.4 PASS-bluff at the recording-integrity layer. On a drop the system MUST automatically (a) DETECT + log an honest offline event (§11.4.6 fabricated-continuity bluff; §11.4.128 always-on gap), (b) WAIT using the project’s ALREADY-DEFINED reconnection timings — never invented (§11.4.6), the SAME grace/reconnect/poll budgets the recovery path already uses, (c) RE-ATTACH + log an honest online/resume event the moment the device returns, (d) ESCALATE to the sanctioned device-recovery path (per §11.4.69 feature class device_recovery) if the device does not return within the defined timeout — through the sanctioned, authorization-gated recovery entry point ONLY, never bypassin its gate, never performing a destructive recovery (power-cycle) autonomously without that authorization (§11.4.21/.101 — high-blast-radius recovery is gated; while blocked keep following + log the blocked-escalation honestly). Composes §11.4.128/.69/.6/.14/.21/.101. Classification: universal (§11.4.17). Prior round: Added §11.4.143 mirror — Real-user-journey mandate for video-streaming-app full-automation tests (User mandate 2026-06-10, verbatim “All video streaming apps ... require to choo some title and to press proper UI button to start or resume playing! Proper UI interaction to play exact show with proper content and subtitles is MANDATORY! Without it we just eventually play on 2nd display sample, and that’s it mostly!”): any full-automation test asserting a video player / streaming app plays content MUST drive the REAL end-user journey through the app’s OWN UI — launch → BROWSE the actual catalog → choose a SPECIFIC title → press the real Play/Resume button → confirm THAT chosen content plays with its correct subtitles on the intended routing target — NEVER a sample / demo / loop-clip / deep-link / am start -a VIEW / synthetic shortcut (those validate ROUTING only, not real chosen-content playback → a §11.4 PASS-bluff at the user-journey layer); non-introspectable streaming UIs use the §11.4.117 CV/OCR pixel oracle; login via the §11.4.10 credential single-source; chosen-content confirmation via §11.4.107 liveness + §11.4.136 real-content + §11.4.137 subtitle correctness; honest operator_attended SKIP-with-reason per §11.4.52/.3 ONLY where autonomous is genuinely infeasible (hard login / geo-block / secure-surface blanking) — NEVER a

Field	Value
	faked routing-only PASS. Composes §11.4.48/.107/.117/.136/.137/.52/.3/§107/§1.1. Classification: universal (§11.4.17). Prior round: Added §11.4.142 mirror — Universal code-review mandate — every change reviewed, always, no exception (User mandate 2026-06-09, verbatim “ALL changes we do MUST pass through the code review step!!! ALWAYS!!!”): EVERY change to ANY governed repository — no exception (source/fixes/tests/gates/mutations/docs/doc-tooling/build-CI/config/governance/conductor-edit sub-agent-output/refactors/one-liners) — MUST pass an INDEPENDENT code-review step BEFORE acceptance/commit/build; the ABSOLUTE form of §11.4.125 (no “just a doc edit”/“just a one-liner”/“self-review §11.4.92 suffices”/“trivial change” carve-out); independence load-bearing (reviewer structurally separated from author per §11.4.70/.20/§11.4.92 self-review PRECEDES never satisfies); anti-bluff (§11.4/.1, conclusions captured-evidence §11.4.5/.69) + iterates to clean GO per §11.4.134; honest boundary §11.4.6 — does NOT replace §11.4.108/.40, it is the FIRST of MULTIPLE STRONG LAYERS. Composes §11.4.1/.4/.5/.6/.20/.40/.69/.70/.92/.108/.110/.125/.134/§107/§1.1. Classification: universal (§11.4.17). Prior round: Added §11.4.132 mirror — Risk-ordered validation priority mandate (User mandate 2026-06-07): tests/validations/verifications MUST run in RISK-DESCENDING order — highest-risk set FIRST (close factors: (a) most-recently-worked, (b) historically most-problematic, (c) highest crash/break/regress likelihood, (d) most-reopened per §11.4.55), each GREEN with real (physical) captured evidence (§11.4.5/.69/.107, no bluff §11.4.6), and ONLY AFTER the set is GREEN does the rest of the suite run; arbitrary-order or lower-risk-before-highest risk-GREEN = §11.4.132 violation; REFINES/STRENGTHENS §11.4.130 + §11.4.46 + §11.4.42 by adding explicit risk-ordering to VALIDATION. Composes §11.4.4/.5/.6/.7/.40/.42/.46/.50/.55/.69/.107/.130. Classification: universal (§11.4.17). Prior round: Added §11.4.131 mirror — Standing session-resumption file mandate (User mandate 2026-06-07): every project MUST maintain a SINGLE canonical always-current session-resumption FILE at a fixed project-declared standard path (declared once per §11.4.35, never moved without a §11.4.66 operator decision) — the OUT-OF-THE-BOX entry point for any fresh session; promotes §11.4.127 (PREPARE-on-demand) into a STANDING version-controlled ARTIFACT, ALWAYS present + in sync, SHORT+FULL variants + live-state anchors + §11.4.65 export + §11.4.44 revision header. Composes §12.10/.127/.65/.44/.6/.66/.126. Classification: universal (§11.4.17). Prior round: Added §11.4.127 mirror — Session-handoff resumption-prompt mandate (User mandate 2026-06-06): when a fresh session is needed (context limits/degradation) OR the operator asks whether a new session is needed, the agent MUST ALWAYS prepare + proactively provide a ready-to-paste resumption prompt valid for that EXACT moment + project phase (SHORT first-sentence + FULL block on demand) pointing to the live handoff doc (.remember/remember.md if present + docs/CONTINUATION.md §12.10, read FIRST git fetch -all), stating PHASE + NEXT action + terminal goal, embedding exact live-state anchors (build IDs/MD5, device serials, HEAD, in-flight PIDs+logs, evidence paths) + binding constraints (anti-bluff §11.4, no-force-push §11.4.113, exact version/naming, hardware gotchas); MUST be moment-valid never generic; handoff docs current BEFORE the prompt; missing/stale/generic = §11.4.127 violation. Composes §12.10/.6/.66/.87/.103/.126. Classification: universal (§11.4.17). Prior round: Added §11.4.126 mirror — Default autonomous-loop working mode from first prompt (User mandate 2026-06-04): the endless fully-autonomous loop is the DEFAULT working mode, engaged automatically the moment the operator sends the FIRST request/prompt of a session — no per-session activation handshake; the loop continues until EITHER (A) a new fully-validated-and-verified version (tag) is created AND published across all owned submodules + main repo to all remotes (per §2.1/§11.4.40/§11.4.113), OR (B) for non-release main-stream work that work is fully completed AND all side work streams done AND nothing is left in the working queue for the current scope; STOPS ONLY on explicit operator STOP, empty current-scope queue, or §12 host-safety; goal ABSOLUTE

Field	Value
	EFFICIENCY; mimicking/imitation/false-results/bluff of ANY kind ABSOLUTELY FORBIDDEN (composes the §11.4 anti-bluff covenant — actually perform the work, never narrate/pretend it); §11.4.126 is the CAPSTONE promoting §11.4.87 from opt-in always-on. Composes §11.4.87/.94/.97/.101/.103/.66/.6/.40/.42/.72/.113 + §2.1 + §12. Classification: universal (§11.4.17). Prior round: Added §11.4.125 mirror — Code-review-agent gate before pre-build + main build (mandatory multi-layer review) (User mandate 2026-06-04): after all batch fixes/changes/implementations are done and BEFORE the pre-build test sweep + main (artifact) build, the agent MUST dispatch dedicated code-review agent(s) (subagent-driven per §11.4.70/.20) that analyze all batch work + all existing data/facts/captured-evidence + the existing codebase (blast radius) + current git history, determine quality + safety + will-it-REALLY-work, and validate+verify every covering test genuinely exercises the work-under-test and catches its negation with ZERO false-result/bluff possibility; any finding (defect / error-prone change / safety risk / will-not-work / bluff-capable test / coverage gap) MUST be fixed+polished+improved+test-covered (four-layer §11.4.4(b), TDD-RED-first §11.4.43/.115) BEFORE the build proceeds; review iterates until no blocking findings; one of MULTIPLE STRONG LAYERS complementing (never replacing) the pre-build sweep + §11.4.92 multi-pass (author-side) + §11.4.108 + §11.4.110. Composes §11.4/.1/.4/.6/.40/.43/.50/.70/.20/.92/.102/.107/.108/.110. Classification: universal (§11.4.17). Prior round: Added §11.4.124 mirror — Dead/unwired-code investigate-before-remove mandate (User mandate 2026-06-04): “zero importers ⇒ dead ⇒ delete” is a guess (§11.4.6), never a finding — before removing ANY seemingly-dead element (zero-importer / never-called / unwired function/type/file/module/package/asset/config/build-target) the agent MUST FIRST investigate via git history (<code>git log --follow, -S/-G</code> pickaxe, blame on the deleted call-site) and capture as FACT where/how it was wired in + when/how it became dead + whether “no references” is real or a hidden reference (reflection/DI/build-tags/codegen/plugin/FFI/config) the static tool cannot see removal is permitted ONLY with captured PROOF it is genuinely unneeded AND MUST be its own separate descriptive commit citing the git-history evidence (+ §11.4.122 operator-confirmation for end-user capabilities; tracked via §11.4.90 Obsolete); with NO such proof the element MUST NOT be removed — instead wire it in properly (§11.4.114) and add any missing/unwired tests (§11.4.27/.43/.115); ALWAYS extra careful — when uncertain, default to NOT removing per §11.4.6/.101/.122. Classification: universal (§11.4.17). Prior round: Added §11.4.123 mirror — Rock-solid proof-or-deep-research mandate (User mandate 2026-06-03): every reported issue / fix / claimed completion MUST be 100% validated with rock-solid CAPTURED proof (§11.4.5/.69/.107) before fixed/implemented/completed; metadata-only / config-only / absence-of-error / grep-without-runtime PASS forbidden (§11.4/.1); when UNSURE how to validate, the agent MUST ALWAYS first perform deep web research (§11.4.8+§11.4.99) to discover/build an evidence-producing method — declaring “untestable” or accepting a metadata-only PASS without exhausting that research is itself a §11.4.123 violation; forensic case study (FACT) 1.1.8-dev FLAG_SECURE + non-introspectable-UI validation unclear → deep research yielded the CV/OCR/liveness/sin probe oracle stack (§11.4.107/.112/.117). Classification: universal (§11.4.17). Prior round: Added §11.4.122 mirror — No-silent-removal-of-existing-components-without-operator-confirmation mandate (User mandate 2026-06-03): no application/component/service/package/feature/driver/module/prebuilt — any already-existing end-user capability — may be removed from the existing codebase/System without FIRST interactively asking the operator (§11.4.66, never silent/autonomous) + an EXPLICIT keep-or-remove decision; silent removal is a release blocker; forensic case study F2 Apple-TV-class app + F4 Huawei HMS component removed without asking, operator reversed both; tracked DROP path ask→approve→Obsolete reason feature-removed+citation (§11.4.90)→remove. Classification: universal (§11.4.17). Prior round

Field	Value
	Added §11.4.114–§11.4.121 mirrors (eight new universal anchors from the 1.1.8-dev live-defect remediation, 2026-06-03): §11.4.114 last-known-good-tag regression isolation; §11.4.115 RED-baseline-on-the-broken-artifact + polarity-switch (refines §11.4.43); §11.4.116 real-time conductor[?]autonomous-test-framework sync channel (append-only event stream + atomically-rewritten status snapshot, verdict events carry evidence paths); §11.4.117 CV/OCR pixel-oracle fallback for non-introspectable UIs (refines §11.4.48/.52/.107); §11.4.118 discovery-pressure to confirm known-issue-set completeness; §11.4.119 single-resource-owner partitioning for parallel hardware testing (refines §11.4.58/.103); §11.4.120 fix-breaks-its-own-gate reconciliation (rewrite the gate, never fake-pass / never revert the fix); §11.4.121 no-commit-while-build-writes-tracked-artifacts (build-output analogue of §11.4.84). All Classification: universal (§11.4.17). §11.4.108–§11.4.113 + §11.4.78–§11.4.107 mirrors continue from earlier Revisions.
Issues	none
Issues summary	—
Fixed	§11.4.108 mirror
Fixed summary	§11.4.107 lands in lockstep with the Constitution.md §11.4.107 addition.
Continuation	—

Table of contents

- [How inheritance works](#)
- [Identity & posture](#)
- [Top-level invariants for every agent](#)
- [Critical base rules restated \(for agents that don't honour @imports\)](#)
 - [Anti-bluff covenant — END-USER QUALITY GUARANTEE \(§11.4\)](#)
 - [Mutation-paired gates \(§1.1\)](#)
 - [Recorded-evidence requirement \(§11.4.2\)](#)
 - [Test-interrupt-on-discovery \(§11.4.4\)](#)
 - [No-guessing mandate \(§11.4.6\)](#)
 - [Item-type tracking \(§11.4.16\)](#)
 - [Universal-vs-project classification \(§11.4.17, User mandate 2026-05-14\)](#)
 - [Subagent-driven-by-default \(§11.4.20, User mandate 2026-05-14\)](#)
 - [Script documentation mandate \(§11.4.18, User mandate 2026-05-14\)](#)
 - [Fixed-document column-alignment \(§11.4.19, User mandate 2026-05-14\)](#)
 - [Operator-blocked status + self-resolution exhaustion \(§11.4.21, User mandate 2026-05-14\)](#)
 - [Document-sync commit discipline \(§11.4.22, User mandate 2026-05-14\)](#)
 - [Build-resource stats tracking \(§11.4.24, User mandate 2026-05-14\)](#)
 - [Full-Automation-Coverage \(§11.4.25, User mandate 2026-05-15\)](#)
 - [Constitution-Submodule Update Workflow \(§11.4.26, User mandate 2026-05-15\)](#)
 - [Submodule-dependency-manifest \(§11.4.31, User mandate 2026-05-15\)](#)
 - [Post-constitution-pull validation \(§11.4.32, User mandate 2026-05-15\)](#)
 - [.gitignore + no-versioned-build-artifacts \(§11.4.30, User mandate 2026-05-15\)](#)
 - [Lowercase-snake_case-naming \(§11.4.29, User mandate 2026-05-15\)](#)

- Submodules-as-equal-codebase + decoupling + dependency-layout (§11.4.28, User mandate 2026-05-15)
- No-fakes-beyond-unit-tests + 100%-test-type-coverage (§11.4.27, User mandate 2026-05-15)
- Type-aware closure-status vocabulary (§11.4.33, User mandate 2026-05-15)
- Reopened-source attribution (§11.4.34, User mandate 2026-05-15)
- Canonical-root inheritance clarity (§11.4.35, User mandate 2026-05-15)
- Mandatory install_upstreams on clone/add (§11.4.36, User mandate 2026-05-15)
- Fetch-before-edit (§11.4.37, User mandate 2026-05-15)
- Installable-asset evidence (§11.4.38, User mandate 2026-05-17)
- Full-suite retest before release tag (§11.4.40, User mandate 2026-05-17)
- Pre-force-push merge-first (§11.4.41, User mandate 2026-05-17)
- Iteration-discipline mandate (§11.4.42, User mandate 2026-05-18)
- TDD-Fix-Discipline (§11.4.43, User mandate 2026-05-18)
- Document revision header (§11.4.44, User mandate 2026-05-18)
- Integration-status-doc maintenance (§11.4.45, User mandate 2026-05-18)
- Validate-recent-work-before-post-flash-tests (§11.4.46, User mandate 2026-05-18)
- Firestore data review (§11.4.47, User mandate 2026-05-18)
- UI-driven video testing (§11.4.48, User mandate 2026-05-18)
- Dual-approach testing (§11.4.49, User mandate 2026-05-18)
- Deterministic consistency (§11.4.50, User mandate 2026-05-18)
- Live-ADB-First maximization (§11.4.51, User mandate 2026-05-18)
- Autonomous-Validation (§11.4.52, User mandate 2026-05-18)
- Fixed_Summary parity (§11.4.53, User mandate 2026-05-18)
- ATM-NNN ticket identifier (§11.4.54, User mandate 2026-05-19)
- Reopens-history tracking + per-item Reopens.md (§11.4.55, User mandate 2026-05-19)
- Status_Summary parity + two-audience format (§11.4.56, User mandate 2026-05-19)
- README.md doc-link section + revision metadata (§11.4.57, User mandate 2026-05-19)
- Parallel-development methodology (§11.4.58, User mandate 2026-05-19)
- README always-sync (§11.4.59, User mandate 2026-05-19)
- Documentation always-sync composite covenant (§11.4.60, User mandate 2026-05-19)
- Credentials-handling (§11.4.10)
- Host-session safety (§12)
- Continuation document (§12.10)
- Data safety (§9)
- CLI workflow expectations
- Hierarchy of project authority
- When stuck

This is the **base AGENTS.md** for any CLI agent (Codex, Cursor, Aider, Continue, Gemini CLI, future LLMs) working on a project that includes the Helix Constitution submodule. Project-level AGENTS.md may extend or tighten rules but MUST NOT weaken them.

Last revision: 2026-05-14

How inheritance works

Because not every CLI agent honours the Markdown `@import` syntax, each consuming project's root `AGENTS.md` MUST do one of:

1. **Reference + restate critical rules** (safest — works for every agent):

```
# <Project> – AGENTS.md

> Base agent rules: `constitution/AGENTS.md` – READ IT FIRST.
> The base file is authoritative for any topic not covered
  here.

## Critical base rules restated (for agents that don't follow
  links)
- Never commit secrets.
- All commits go through the project's commit wrapper.
- Anti-bluff covenant binds every test (Constitution §11.4).

## Project-specific
- ...
```

2. **Generate at build time** — keep one source-of-truth in this submodule, concatenate with a project-specific delta file. A `compose-agents.sh` script in this submodule's `Upstreams/` (or `scripts/` if added) would output `constitution/AGENTS.md` followed by `---` followed by `AGENTS.project.md`.

Identity & posture

You are working inside a project that has opted into the Helix Constitution. The Constitution is the source of truth for engineering discipline (test coverage, anti-bluff covenant, data safety, host safety, credentials handling, documentation discipline). Whenever this `AGENTS.md` mentions a `§X.Y` reference, it points to `constitution/Constitution.md`.

The Constitution applies recursively to every submodule of every project that includes it — you **MUST** treat the Constitution's rules as in-force regardless of how deep into the submodule tree you are working.

Top-level invariants for every agent

1. **Never assume; verify by reading files.** Especially before any destructive action.
2. **Never bluff.** Do not invent file paths, command outputs, test results, or library APIs you have not verified. If you are unsure, say so explicitly and verify.
3. **No guessing language** (§11.4.6). `likely`, `probably`, `maybe`, `might`, `appears`, `seems` are forbidden when reporting causes. Use captured evidence or mark explicitly `UNCONFIRMED`.
4. **Test coverage for every change** (§1) — pre-build gate, post-build gate, runtime test, AND a meta-test mutation that proves the gate catches regressions.

5. **Commit through the project's official wrapper.** Direct `git add / git commit / git push` on the main repo is forbidden in normal workflow.
6. **Never skip hooks** (`--no-verify`, `--no-gpg-sign`).
7. **Never force-push.** Force-push requires explicit per-session human authorization AND a green §9.1.5 post-op gate.
8. **Hardlinked backup before any destructive op** (§9). Zero excuse.
9. **Host safety** (§12) — abort if pre-flight check fails; wrap heavy work in bounded execution scopes; never exceed 60% of host RAM.
10. **Credentials NEVER tracked** (§11.4.10) — `.env` patterns git-ignored; runtime-load only; per-service file separation. **Pre-store leak audit** (§11.4.10.A, 2026-05-17) — before storing operator-provided credentials in any gitignored config, grep the entire tracked tree AND git history for the literal value(s); surface any findings to operator before storing.
11. **CONTINUATION.md kept in sync** (§12.10) — every non-trivial state change updates the continuation document in the same commit.

Critical base rules restated (for agents that don't honour @imports)

Anti-bluff covenant — END-USER QUALITY GUARANTEE (§11.4)

Forensic anchor — verbatim user mandate (2026-04-28):

“We had been in position that all tests do execute with success and all Challenges as well, but in reality the most of the features does not work and can't be used! This MUST NOT be the case and execution of tests and Challenges MUST guarantee the quality, the completion and full usability by end users of the product!”

The bar for shipping is **not** “tests pass” but “**users can use the feature.**” Every PASS MUST carry positive evidence captured during execution that the feature works for the end user. Metadata-only PASS, configuration-only PASS, absence-of-error PASS, and grep-based PASS without runtime evidence are critical defects.

Mutation-paired gates (§1.1)

Every new gate MUST be paired with a mutation entry in the project's meta-test harness that mutates the asserted condition and asserts the gate now FAILs. A gate without a paired mutation is a bluff gate and is forbidden.

Recorded-evidence requirement (§11.4.2)

Every PASS for a user-visible feature MUST be cross-checked against a captured recording + action timeline. A PASS that lacks at least one matched timeline event in the analyzer findings is a §11.4 PASS-bluff.

Test-interrupt-on-discovery (§11.4.4)

The moment any defect is re-discovered, re-produced, or newly identified during a test cycle, the cycle MUST stop. Then: systematic debugging → fix at root cause → four-layer test coverage (pre-build / post-build / runtime / meta-test paired mutation) → full rebuild → re-deploy on every target → full retest from beginning.

No-guessing mandate (§11.4.6)

Forensic anchor — verbatim user mandate (2026-05-08):

“‘LIKELY’ is guessing, we MUST NOT have guessing, since it can be or may not be! No bluffing and uncertainty is allowed at any cost! We MUST always know exactly precisely what is happening exactly, in any context, under any conditions, everywhere!”

Forbidden vocabulary when reporting causes: likely, probably, maybe, might, possibly, presumably, seems, appears to, guess, seemingly, apparently, perhaps, supposedly, conjectured, and synonyms.

Item-type tracking (§11.4.16)

Every active item in the project Issues file MUST carry a ****Type:**** line within eight non-blank lines of its heading. Three-value CLOSED vocabulary: Bug (product defect / regression / user-visible broken behaviour), Feature (new capability not previously offered to end users), Task (internal workstream — refactor, doc, infra, gate, audit; the lowest-stakes default when ambiguous). The Issues_Summary file carries the Type column for every active item. All three Issues / Issues_Summary / Fixed file types kept in sync (Markdown + HTML + PDF). Pre-build gates CM-ITEM-TYPE-TRACKING + CM-COVENANT-114-16-PROPAGATION enforce the mandate. No escape hatch.

Universal-vs-project classification (§11.4.17, User mandate 2026-05-14)

Before adding ANY new rule / mandatory constraint / covenant clause / gate / “MUST”-bearing statement, classify it: **universal** (any project → constitution submodule) or **project-specific** (specific hardware / vendor / package / region → project / submodule layer). Commit message MUST carry **Classification:** line + one-sentence rationale. When uncertain, default to project-specific. Gate CM-UNIVERSAL-VS-PROJECT-CLASSIFICATION audits new rule commits. Paired mutation enforces the gate is not a bluff.

Subagent-driven-by-default (§11.4.20, User mandate 2026-05-14)

When the runtime supports subagent delegation (Agent tool / task runner / sub-session), the primary agent MUST default to subagent delegation for multi-step scope (≥ 3 phases), parallelisable independent subtasks, long-running diagnostic loops, OR specialised domain workflows. Foreground-only is reserved for single-file edits, mid-execution operator clarification, critical-state sequencing (commits / pushes / tags), or tasks under ~30s. Discipline: tight scope (4-6 tasks), checkpoint

commits per task, anti-stall protection in prompts, anti-bluff verification of subagent claims via repo state. Parallel subagents MUST partition non-overlapping files; parent commit_all.sh --auto-cascade bundles via git add -A. Gate CM-SUBAGENT-DELEGATION-AUDIT (per consuming project) flags foreground multi-step work as a violation. Paired mutation enforces the gate is not a bluff.

Script documentation mandate (§11.4.18, User mandate 2026-05-14)

Every Bash / shell / POSIX-sh script under scripts/ / bin/ / tests/ / library dirs / CI hooks (depth-N recursive) MUST carry (1) an in-source documentation block at the top (Purpose / Usage / Inputs / Outputs / Side-effects / Dependencies / Cross-references) AND (2) an external user guide under docs/scripts/<script-name>.md (Overview / Prerequisites / Usage examples / Edge cases / Internal behaviour / Related scripts / Last verified date). When a script is modified, BOTH the in-source block AND the external user guide MUST be updated in the SAME commit. **No documentation ever out of sync with its codebase.** Gate CM-SCRIPT-DOCS-SYNC enforces. Paired mutation strips the invariant.

Fixed-document column-alignment (§11.4.19, User mandate 2026-05-14)

The closed-archive tracker (Fixed.md) MUST mirror the open-work tracker (Issues.md + Issues_Summary.md) along the same lifecycle axes — at minimum **Status** and **Type**. Concretely: (1) every #### / ##### heading in Fixed.md carries ****Status:**** (from closed-set {Fixed (→ Fixed.md) | Fixed – pending device verification | Fixed – RECLASSIFIED}) and ****Type:**** ({Bug | Feature | Task}) within 8 non-blank lines of the heading; (2) Fixed_Summary.md companion exists with column structure # | Level | Status | Type | One-line description matching Issues_Summary.md exactly; (3) all three formats (.md, .html, .pdf) for BOTH Fixed and Fixed_Summary stay in sync via the same single-shot wrapper that handles Issues + Issues_Summary; (4) closure is atomic — resolved items migrate from Issues.md to Fixed.md in the same commit. Pre-build gate CM-FIXED-COLUMN-ALIGNMENT (5+ invariants: file-exists, header-has-Status+Type, mtime ≥ Fixed.md, generator script present, sync wrapper invokes it, HTML+PDF exports present). Paired mutation strips Status column from Fixed_Summary header → gate FAILs. Classification: universal (per §11.4.17).

Operator-blocked status + self-resolution exhaustion (§11.4.21, User mandate 2026-05-14)

Operator-blocked is the §11.4.15 Status closed-set's 7th value alongside {Queued | In progress | Ready for testing | In testing | Reopened | Fixed (→ Fixed.md)}. **Last-resort classification** — agent MUST first exhaust applicable self-resolution paths: (a) CLI / ADB / SSH / API access already available, (b) subagent delegation per §11.4.20, (c) existing repo tooling, (d) captured fallback (synthetic event / asset substitution / §11.4.3 SKIP), (e) external research per §11.4.8. Every Operator-blocked item MUST carry ****Operator-Block-Details:**** within 8 non-blank lines of its heading stating WHAT (action) / WHY (each exhausted alternative) / UNBLOCK CONDITION (observable signal) / WHO (contact or doc pointer). Issues_Summary.md lists it as a sortable Status value. Items re-evaluated every Nth tag cycle (≥3rd recommended). Fake Operator-

blocked without exhaustion audit = §11.4 covenant violation at planning layer (PASS-bluff equivalent). Gates CM-ITEM-OPERATOR-BLOCKED-DETAILS + CM-OPERATOR-BLOCKED-SELF- RESOLUTION-AUDIT (every NEW Operator-blocked commit carries an “Attempted: a — ...; b — ...; c — ...” trail). Classification: universal (per §11.4.17). No escape hatch.

Document-sync commit discipline (§11.4.22, User mandate 2026-05-14)

Every project tracking work items through an Issues / Fixed lifecycle MUST provide a **lightweight commit path** distinct from the full-repo commit wrapper. Stages, commits, pushes ONLY the status-tracking doc set: Issues + Issues_Summary + Fixed + Fixed_Summary + CONTINUATION + their HTML + PDF exports + auto-generated audit artifact. Wrapper MUST: (a) auto-invoke the project’s export-regeneration pipeline first so Markdown + HTML + PDF stay in sync; (b) stage ONLY the explicit doc-set list — NEVER `git add -A`; (c) use a separate flock disjoint from the full-tree wrapper; (d) push to every parent-repo remote; (e) exit 3 on nothing-to-commit (informational). MUST be invocable standalone OR via a delegation flag on the full-tree wrapper (e.g. `commit_all.sh --docs-only`). Inherits §9 preflight. Gate CM-COMMIT-DOCS-EXISTS + paired mutation. Composes with §11.4.12 / §11.4.15 / §11.4.18 / §12.10. Classification: universal (per §11.4.17). No escape hatch — doc-status drift is a §11.4 PASS-bluff at the documentation layer.

Build-resource stats tracking (§11.4.24, User mandate 2026-05-14)

Every project under this Constitution with a build exceeding 1 minute wall-clock MUST run a host-side resource sampler for every build — samples `/proc/meminfo` + `/proc/loadavg` + `/proc/stat` + `/proc/diskstats` at a fixed interval (recommended 5 s); writes TSV. On stop, computes **min / max / mean / p95** per metric; appends one TSV row per build to a registry; regenerates a Markdown report (Stats.md) whose top surfaces **ever-values** (min / max / mean across all tracked builds) and whose body lists per-build entries most-recent-first with SUCCESS / FAIL / UNKNOWN + reason. Stats.md exported to Stats.html + Stats.pdf through the project’s normal export pipeline per §11.4.12. Triple committed via the §11.4.22 lightweight doc-sync wrapper. Sampler MUST stay under 50 MB RSS / 5% CPU. Stop hook called from both success AND failure paths of the build wrapper. Gate CM-BUILD-RESOURCE-STATS-TRACKER + paired mutation. Classification: mixed (per §11.4.17). Composes with §11.4.12 / §11.4.18 / §11.4.22 / §12.6 / §12.7 / §12.9. No escape hatch — build-resource debugging without time-series data is the bluff this anchor forbids.

Full-Automation-Coverage (§11.4.25, User mandate 2026-05-15)

Forensic anchor — verbatim user mandate (2026-05-15):

“Make sure that every feature, every functionality, every flow, every use case, every edge case, every service or application, on every platform we support is covered with full automation tests which will confirm anti-bluff policy and provide the proof of fully working capabilities, working implementation as expected, no issues, no bugs, fully documented, tests covered!”

No feature / functionality / flow / use case / edge case / service / application on any supported platform may be considered deliverable until automation tests prove six invariants: (1) anti-bluff posture with captured runtime evidence (§7.1 + §11.4); (2) proof of working capability end-to-end on target topology (§11.4.3, no mocks); (3) implementation matches documented promise; (4) no open issues/bugs surfaced (cross-checked vs §11.4.15 / §11.4.16); (5) full documentation (user manual + §11.4.18 for scripts) kept in sync via §11.4.12; (6) four-layer test floor per §1 (pre-build + post-build + runtime + paired mutation). Consuming projects publish a coverage ledger (feature × platform × invariant × status) regenerated at release-gate sweep. Classification: universal (§11.4.17). Severity-equivalent to a §11.4 PASS-bluff at the release-gate layer. No escape hatch. See Constitution §11.4.25 for the full mandate.

Constitution-Submodule Update Workflow (§11.4.26, User mandate 2026-05-15)

Forensic anchor — verbatim user mandate (2026-05-15):

“Every time we add something into our root (constitution Submodule) Constitution, CLAUDE.MD and AGENTS.MD we MUST FIRST fetch and pull all new changes / work from constitution Submodule first! All changes we apply MUST BE committed and pushed to all constitution Submodule upstreams! In case of conflict, IT MUST BE carefully resolved! Nothing can be broken, made faulty, corrupted or unusable! After merging full validation and verification MUST BE done!”

Before any modification to constitution/{Constitution, CLAUDE, AGENTS}.md, execute in order: (1) fetch + pull (–ff-only or –rebase; never strategy=ours / unrelated-histories without authorization) so the submodule is at upstream tip before editing; (2) apply the change with §11.4.17 classification + verbatim mandate quote; (3) validate (meta_test_inheritance.sh + no merge-conflict markers + cross-file consistency); (4) commit (governance files only, no git add –A) + push to EVERY configured upstream remote (§2.1 violation if not); (5) careful conflict resolution preserving union of governance content — force-push forbidden (§9.2); (6) post-merge validation: git submodule update –remote –init + re-run cascade verifier (CONST-047) confirming the new clause reaches every owned submodule; (7) bump consuming project’s .gitmodules pointer to new HEAD in the SAME commit as cascade work — out-of-sync pointer = §11.4.26 violation. Classification: universal (§11.4.17). Severity-equivalent to force-push without §9.2 authorization. No escape hatch. See Constitution §11.4.26 for the full pipeline (operational scope, cross-cutting reach).

Submodule-dependency-manifest (§11.4.31, User mandate 2026-05-15)

Every owned-by-us submodule MUST ship helix-deps.yaml listing its own-org dependencies: {name, ssh_url, ref, why, layout: flat|grouped}. Tooling incorporate-submodule <ssh-url> adds the submodule at the parent project’s canonical path (CONST-051(C)) + reads its helix-deps.yaml + recurses + aborts on conflicting refs + emits <root>/helix-manifest.yaml audit record. Anti-bluff: each manifest paired with a Challenge that bootstraps from scratch, asserts layout matches manifest, runs submodule tests, captures wire evidence. Without the proof, the manifest is a §11.4.31 violation. This rule is the operational complement of

§11.4.28 / CONST-051(C) — manifests are the bridge that lets consumers reconstruct the dependency graph at the parent root. Classification: universal (§11.4.17). See Constitution §11.4.31 for the full mandate.

Post-constitution-pull validation (§11.4.32, User mandate 2026-05-15)

Whenever a project's constitution submodule is fetched + pulled with any content change, run `scripts/verify-all-constitution-rules.sh` BEFORE the new HEAD is treated as canonical for other work. The sweep re-runs the governance-cascade verifier + every implementable rule gate (CONST-053 .gitignore audit, CONST-051(C) nested-own-org audit, CONST-052 case audit, CONST-050(A) mock-from-production audit, CONST-035 anti-bluff smoke). Failures populate Issues per §11.4.15 (Status: Reopened, Type: Bug); closure requires positive-evidence per §11.4. Pull-time invocation auto-triggered by `git submodule update --remote constitution`. Anti-bluff: sweep's own meta-test (paired mutation §1.1) plants a violation per gate, asserts sweep FAILs. This is the **enforcement engine** for every other rule — without it, new rules are decorative anchors rather than enforced gates. Classification: universal (§11.4.17). See Constitution §11.4.32 for the full mandate.

.gitignore + no-versioned-build-artifacts (§11.4.30, User mandate 2026-05-15)

Forensic anchor — verbatim user mandate (2026-05-15):

“every project module, every Submodule, every service and application MUST HAVE proper .gitignore file! We MUST NOT git version build artifacts, cache files, tmp files, main .env file(s) or any files containing sensitive data, API keys or token! ... If any violation is detected it MUST be fixed before commit is executed!”

Every project module / submodule / service / application MUST ship a proper .gitignore. Forbidden-from-version-control: build artefacts (`bin/`, `build/`, `dist/`, `target/`, `*.exe`, `*.so`, `*.class`, `*.pyc`), cache files (`__pycache__/`, `node_modules/`, `.gradle/`, `.terraform/`), temp files (`*.tmp`, `*.swp`, `.DS_Store`), sensitive data (`.env`, `.env.*` except `.env.example` placeholder, `*.pem`, `*.key`, `id_rsa*`, `.netrc`, `secrets/`, `api_keys.sh`), generated logs (`*.log`, `coverage.out`, `htmlcov/`), OS/IDE personal state (`.idea/`, `.history/`).

Anti-bluff invariant: ignore-line alone is not enough — no file matching forbidden patterns may be tracked. Pre-commit attention: every author inspects `git diff --staged + git status` BEFORE commit; violations abort the commit (un-stage, add to ignore, scrub if already-tracked). Gate `CM-GITIGNORE-PRECOMMIT-AUDIT` + paired mutation. Recreatable-content test: if a documented mechanism regenerates the file from sources, it's a build derivative and MUST be ignored.

Secret-leak intersection (§11.4.10 / CONST-042 / §12.1): a .env leak is BOTH a §11.4.30 and a §11.4.10 violation; rotation + post-mortem required.

Classification: universal (§11.4.17). Severity-equivalent to §11.4 PASS-bluff at the repository-hygiene layer. Composes with §1, §2, §9.1, §11.4.10, §11.4.12, §11.4.17, §11.4.18, §11.4.20, §11.4.25, §11.4.26, §11.4.27, §11.4.28, §11.4.29, CONST-047. See Constitution §11.4.30 for the full mandate.

Lowercase-snake_case-naming (§11.4.29, User mandate 2026-05-15)

Forensic anchor — verbatim user mandate (2026-05-15):

“naming convention for Submodules and directories (applied deep into hierarchy recursively) - all directories and Submodules MSUT HAVE lowercase names with space separator between the words of ‘_’ character (snake-case)! All existing Submodules and directories which are not following this rule MUST BE renamed! ... NOTE: Rules lowercase / snake-case do apply to all project files as well and references to it and from them!”

Every directory / submodule / file MUST use lowercase snake_case. Non-compliant names MUST be renamed; every reference (configs, docs, source imports, governance files) updated atomically (reference drift = §11.4.29 violation of equal severity). Exceptions (common-sense, technology-preserving): language- mandated case for Java/Kotlin/Android/Apple/C#/Swift inside the language root; vendor third-party submodules; build artefacts.

Upstreams/ → upstreams/ transition: constitution submodule’s install_upstreams.sh MUST support BOTH directory names during migration; lowercase wins when both present.

Project-Toolkit Upstreamable machinery fetched+pulled before any rename batch + itself complies. Lacking BOTH-dir support is a release blocker.

Test coverage of renames: regression test for reference resolution + full CONST-050(B) test-type matrix + anti-bluff wire-evidence.

Phased execution: brainstorming → phase plan → fine-grained tasks/subtasks → every change covered. §11.4.20 subagent delegation for cross-cutting rename sweeps.

Classification: universal (§11.4.17). No escape hatch beyond common-sense exceptions. Severity-equivalent to §11.4 PASS-bluff at the reference-integrity layer. Composes with §1, §11.4.12, §11.4.17, §11.4.18, §11.4.20, §11.4.25, §11.4.26, §11.4.27, §11.4.28, CONST-047. See Constitution §11.4.29 for the full mandate.

Submodules-as-equal-codebase + decoupling + dependency-layout (§11.4.28, User mandate 2026-05-15)

Forensic anchor — verbatim user mandate (2026-05-15):

“All existing Submodules in the project that we are controlling and belong to some our organizations (vasic-digital, HelixDevelopment, red-elf, ATMOSphere1234321, Bear-Suite, BoatOS123456, Helix-Flow, Helix-Track, Server-Factory — we can ALWAYS check dynamically using GitHub and GitLab CLIs) are equal parts of the project’s codebase! We MUST work on that code as much as we do with main project’s codebase! ... We MUST NEVER modify Submodules to bring into them any project specific context

... All Submodule dependencies that are used by Submodule MUST BE accessed from the root of the project! We MUST NOT have nested Submodule dependencies.”

Three invariants. **(A) Equal-codebase:** every owned-by-us submodule (orgs: vasic-digital, HelixDevelopment, red-elf, ATMOSphere1234321, Bear-Suite, BoatOS123456, Helix-Flow, Helix-Track, Server-Factory — discoverable via gh/ glab) is an equal part of the consuming project’s codebase. Same engineering attention: analysis, extension, tests, gap-fill, bug-fix, documentation. Coverage ledgers list each submodule as in-scope. **(B) Decoupling:** NEVER inject project-specific context INTO submodules; they remain project-not-aware, reusable, modular, testable. When a submodule needs parent info, use configuration injection. **(C) Dependency-layout:** every dependency consumed by an owned submodule lives at <root>/<name>/ or <root>/submodules/<name>/ of the parent project. Nested own-org submodule chains FORBIDDEN — add the dependency at the parent root; the submodule reaches it via documented import/SDK/runtime resolver. Third-party submodules exempt.

Gates: CM-OWNED-SUBMODULE-EQUAL-ENGINEERING, CM-OWNED-SUBMODULE-DECOUPLING, CM-OWNED-SUBMODULE-LAYOUT. Paired mutations for each. Classification: universal (§11.4.17). No escape hatch. Severity-equivalent to a §11.4 PASS-bluff at the codebase- completeness layer. Composes with §1, §3, §11.4.17, §11.4.20, §11.4.25, §11.4.26, §11.4.27, CONST-047. See Constitution §11.4.28 for the full mandate.

No-fakes-beyond-unit-tests + 100%-test-type-coverage (§11.4.27, User mandate 2026-05-15)

Forensic anchor — verbatim user mandate (2026-05-15):

“Mocks, stubs, placeholders, TODOs or FIXMEs are allowed to exist ONLY in Unit tests! All other test types MUST interact with real fully implemented System! No fakes, empty implementations or bluffing is allowed of any kind! All codebase of the project MUST BE 100% covered with every supported test type.”

Two invariants. **(A)** Mocks/stubs/fakes/placeholders/TODOs/ FIXMEs/“for now” patterns PERMITTED only in unit-test sources; non-unit tests (integration, E2E, full-automation, security, DDoS, scaling, chaos, stress, performance, benchmarking, UI, UX, Challenges, HelixQA) MUST exercise the real, fully implemented system. Production code MUST NOT import mock paths. Gate CM-NO-FAKES-BEYOND-UNIT-TESTS + paired mutation. **(B)** Codebase MUST be covered by every supported test type the domain warrants: unit, integration, E2E, full-automation, security, DDoS, scaling, chaos, stress, performance, benchmarking, UI, UX, Challenges (vasic-digital/Challenges submodule fully incorporated), HelixQA (HelixDevelopment/HelixQA submodule fully incorporated, with full autonomous QA sessions executing every registered test bank).

Required dependency submodules (recursive per CONST-047): Challenges (git@github.com:vasic-digital/Challenges.git) + HelixQA (git@github.com:HelixDevelopment/HelixQA.git) + any other functionality submodules under vasic-digital/HelixDevelopment orgs the project depends on. Pointers bumped to upstream HEAD in same commit as cascade work (§11.4.26 step 7).

Classification: universal (§11.4.17). Severity-equivalent to a §11.4 PASS-bluff at the release-gate layer. No escape hatch. Composes with §1, §7.1, §11.4.1–§11.4.26 (esp. §11.4.25 — strict per-type-of-test expansion). See Constitution §11.4.27 for full mandate.

Type-aware closure-status vocabulary (§11.4.33, User mandate 2026-05-15)

Every project that tracks work items by Type per §11.4.16 **MUST** close them with the Type-appropriate closure-status word: Bug → Fixed (→ Fixed.md), Feature → Implemented (→ Fixed.md), Task → Completed (→ Fixed.md). The (→ Fixed.md) suffix is preserved across all three so existing migration tooling (atomic Issues.md → Fixed.md move per §11.4.19) keeps working. Generators treat the three terminal values as semantically equivalent (all closed, positive evidence captured) but preserve the literal in emitted docs. Closing a Feature with Fixed (→ Fixed.md) or a Task with Implemented (→ Fixed.md) is a §11.4.33 violation. Pre-build gate CM-CLOSURE-VOCAB-TYPE-AWARE. Composes with §11.4.15 / §11.4.16 / §11.4.19 / §11.4.23. Classification: universal (per §11.4.17). No escape hatch.

Reopened-source attribution (§11.4.34, User mandate 2026-05-15)

Every Issues.md heading whose ****Status:**** is Reopened **MUST** carry a ****Reopened-Details:**** line within 8 non-blank lines of the heading, capturing four sub-facts: **By:** AI or User; **On:** ISO date; **Reason:** one of { test-failed | manual-testing-detected | captured-evidence-contradicts | end-user-report | cycle-re-discovered | design-reconsidered } or explicit free text; **Evidence:** path or short description of the captured artefact. Reopens without evidence are §11.4.6 / §11.4.7 violations: the reopen IS a demotion-from-Fixed change. Issues_Summary.md Status column **MUST** distinguish Reopened sub-states by source (e.g., Reopened (AI: test-failed) vs Reopened (User: manual-testing)). Pre-build gate CM-ITEM-REOPENED-DETAILS mirrors §11.4.21 walk pattern. Composes with §11.4.6 / §11.4.7 / §11.4.15 / §11.4.21. Classification: universal (per §11.4.17). No escape hatch.

Canonical-root inheritance clarity (§11.4.35, User mandate 2026-05-15)

The constitution submodule's three files (constitution/Constitution.md, constitution/CLAUDE.md, constitution/AGENTS.md) ARE the canonical root — also called the parent files. Universal rules per §11.4.17 live here.

The consuming project's repository-root files (<project-root>/CLAUDE.md, <project-root>/AGENTS.md, optionally <project-root>/Constitution.md or equivalent) are consumer extensions. They open with the inheritance pointer (either @constitution/CLAUDE.md import or ## INHERITED FROM constitution/CLAUDE.md heading). Project-specific rules per §11.4.17 live here.

When in doubt: universal rule → constitution submodule; project-specific rule → consumer's file. Default consumer-side when uncertain. “Parent CLAUDE.md” / “root Constitution” → constitution submodule file at constitution/<filename>. “Project CLAUDE.md” / “this project's AGENTS.md” → consumer file at

<project-root>/<filename>. Moving a rule between layers MUST be a visible commit — `git mv` + explicit “Lifted from to constitution per §11.4.35” / “Demoted from constitution to per §11.4.35” line in the message. Pre-build gate CM-CANONICAL-ROOT-CLARITY. Composes with §11.4.17. Classification: universal (per §11.4.17). No escape hatch.

Mandatory install_upstreams on clone/add (§11.4.36, User mandate 2026-05-15)

Forensic anchor — verbatim user mandate (2026-05-15):

“Every Submodule or Git repository we add or clone MUST BE upstreams installed using Upstreamable utility which MUST BE available through exported paths of the host system (in `.bashrc` or `.zshrc`) using `install_upstreams` command executed from the root of the cloned (added) repository - only if in it is Upstreams or upstreams directory present with bash script files (recipes) for all repository’s upstreams!”

Every clone / add of a Git repository under any consuming project MUST be followed by `install_upstreams` invocation from the repository’s root IF its tree contains `upstreams/` (or legacy `Upstreams/` per §11.4.29 transition) populated with `*.sh` recipe files. The utility (installed on operator’s PATH via `.bashrc/ .zshrc`, implementation lives in this constitution submodule) reads recipe files + configures every declared upstream as a named git remote + fans out origin push URLs.

Skipping the invocation when `upstreams/` is present silently breaks §2.1 (multi-upstream push is the norm). Gate CM-INSTALL-UPSTREAMS-ON-CLONE + paired mutation. Automation: `incorporate-submodule` (§11.4.31) auto-invokes; manual invocation also supported. Pre-commit check: `git remote -v | grep -c push` reports expected upstream count.

Classification: universal (§11.4.17). Composes with §2, §2.1, §3, §9.2, §11.4.17, §11.4.20, §11.4.28, §11.4.29, §11.4.30, §11.4.31. See Constitution §11.4.36 for the full mandate.

Fetch-before-edit (§11.4.37, User mandate 2026-05-15)

Forensic anchor — verbatim user mandate (2026-05-15):

“Make sure that `feedback_fetch_before_edit` memory rule is part of our constitution Submodule - the root Consitution, AGENTS.MD and CLAUDE.MD. Validate and verify that Proejct-Toolkit and all Submodules do inherit all of them!”

FIRST git-touching action of any session: `git fetch --all --prune && git log --oneline HEAD..@{u} && git submodule foreach --recursive 'git fetch --all --prune --quiet'`. If `HEAD..@{u}` is non-empty, integrate (ff-merge / rebase / surface per §11.4.4) BEFORE any local edit, scanner, or test. Skipping on “nothing could have changed” is a §11.4.6 (no-guessing) violation — remote state is not knowable without a fetch.

The originating 2026-05-15 incident: agent ran `git remote remove gitee` as a fresh task, but the work had been committed upstream by a parallel agent 25 minutes earlier — local config edit was a no-op echo of already-merged history. A 30-second fetch would have caught it.

Scope: consuming project root + every owned submodule recursively (§11.4.28) + the constitution submodule itself (§11.4.26 step 1) + any dependency cloned via `incorporate-submodule` (§11.4.31) or `git submodule add` (§11.4.36). Gate CM-FETCH-BEFORE-EDIT-AUDIT + paired mutation per §1.1. Classification: universal (per §11.4.17). No escape hatch.

Installable-asset evidence (§11.4.38, User mandate 2026-05-17)

For any user-distributable build artifact (package, bundle, installer, or container image), tests/challenges MUST open the artifact and verify each user-visible asset is **present** and **non-degenerate**. A PASS without artifact-opening verification is a §11.4 PASS-bluff. The failure mode: source file exists → build packages it → source-layer checks pass → artifact produced with asset stripped or misconfigured, and no gate opens the artifact to verify.

Required per consuming project: one challenge script per artifact type that opens the produced artifact and verifies every declared user-visible asset. The challenge MUST run as part of the standard QA gate. Classification: universal (per §11.4.17). No escape hatch. See Constitution §11.4.38 for the full mandate.

Full-suite retest before release tag (§11.4.40, User mandate 2026-05-17)

Release tag MUST NOT be created until a COMPLETE retest with ALL existing tests runs on a clean baseline AFTER every workable item in the batch is done, fixed, polished, individually verified. Spot-check retests FORBIDDEN — miss interaction defects. Comprises 7 steps: (1) pre-build full sweep, (2) post-build full sweep, (3) on-device 4-phase cycle on EVERY owned device, (4) meta-test full mutation sweep, (5) Challenge bank full sweep, (6) Issues.md/Fixed.md state audit, (7) CONTINUATION.md sync check. Typically 12–48h elapsed, NOT optional, NOT abbreviated. Composes with §11.4.4 (per-fix retest still required at fix granularity) + §11.4.7 (full-suite is authoritative baseline for closures) + §11.4.39 (per-feature on-device validation runs as step 3). Classification: universal (per §11.4.17). No escape hatch. See Constitution §11.4.40 for the full mandate.

Pre-force-push merge-first (§11.4.41, User mandate 2026-05-17)

Forensic anchor — verbatim user mandate (2026-05-17):

“make sure we bring everything from branches to our side before forc push is done! Afer everything is safely and fully merged and all potential conflicts (if any) resolved, then do force push! make sure nothing isnlost, broken or corrupted on bith sides!”

Any force-push authorised under CONST-043 MUST be preceded by a mechanical 4-step merge-first pipeline: (1) fetch every remote (`git fetch --all --prune --tags`); (2) integrate every divergent commit locally (`rebase / merge / operator-`

confirmed cherry-pick per appropriate strategy); (3) audit the integrated tree (no conflict markers anywhere in governance + source + test files; no file silently dropped; previously-passing tests still pass; captured-evidence artefacts still validate); (4) only THEN execute `git push --force-with-lease`.

Two-gate composition with CONST-043. §11.4.41 does NOT relax CONST-043 — it adds a SECOND mechanical gate. CONST-043 alone authorises a push that loses remote work; §11.4.41 alone risks pushing without operator awareness. Both required.

Three failure modes prevented: remote-side content loss (parallel sessions overwritten), stale-state acts (`--force-with-lease` reads stale local refs without prior fetch), conflict-driven corruption (markers committed verbatim — observed 2026-05-17 in `helix_qa` + containers governance files). Verification artefact: `docs/changelogs/<tag>.md` “Force-push merge-first audit” section captures fetch output, divergence log, integration strategy, conflict-marker scan, test delta, push output with lease SHA, + CONST-043 authorisation quote. Gate `CM-FORCE-PUSH-MERGE-FIRST` + paired mutation. Classification: universal (§11.4.17). No escape hatch. Composes with §9.2, §11.4.4, §11.4.6, §11.4.26, §11.4.32, §11.4.37, §11.4.40, CONST-043, CONST-047. See Constitution §11.4.41 for the full mandate. ### Iteration-discipline mandate (§11.4.42, User mandate 2026-05-18)

Work proceeds in priority-ordered cycles. Five mandatory steps per cycle: (1) select TOP + MIDDLE critical only (defer LOW until critical batch closed), (2) batch implementation per §11.4.4 + §11.4.9, (3) smoke gate (<30 min), (4) ONLY if smoke GREEN and no new user/operator report → §11.4.40 full retest (12–48 h), (5) release-ready OR loop back. Composes with §11.4.4 / §11.4.7 / §11.4.9 / §11.4.34 / §11.4.40 — §11.4.42 is the meta-loop conductor binding them. Every PASS in smoke and full retest MUST carry positive captured evidence (§11.4.2 + §11.4.5). Tests AND Challenges bound equally. No escape hatch. Classification: universal (per §11.4.17). See Constitution §11.4.42 for the full mandate.

TDD-Fix-Discipline (§11.4.43, User mandate 2026-05-18)

Every fix follows 5-step workflow: **RED** (failing test FIRST, real product defect, captured evidence) → **LIVE-ADB-PROBE** (try on running device when feasible; INFEASIBLE for kernel / framework / HAL / vendor / `init.rc` / `sepolicy` / `ro.*` / `Android.bp` — must rebuild; cite `LIVE_PROBE_INFEASIBLE: <reason>` in commit) → **GREEN** (source patch, batched per §11.4.9, four-layer coverage per §11.4.4) → **VERIFY** (re-run RED test, PASSes under SAME conditions per §11.4.7, 10-iteration reliability per §11.4.42) → **DOCUMENT** (Issues.md → Fixed.md with type-aware closure §11.4.33, CLAUDE.md Applied Fixes row, changelog, guides, HelixQA bank, CONTINUATION.md §12.10 — all in SAME commit). Test-after-fix is a §11.4 PASS-bluff: test agrees with fix but does not catch the bug. No escape hatch. Classification: universal (§11.4.17). See Constitution §11.4.43 for the full mandate.

Document revision header (§11.4.44, User mandate 2026-05-18)

Every IN-scope tracked Markdown document (Issues.md / Issues_Summary.md / Fixed.md / Fixed_Summary.md / CONTINUATION.md / docs/guides/** / docs/research/** / docs/scripts/** / docs/changelogs/** / docs/superpowers/plans/** / docs/hardware/** / every other docs/*.md) carries `**Revision:** N` (monotonic

positive integer) + ****Last modified:**** YYYY-MM-DDTHH:MM:SSZ directly below the H1 title. CLAUDE.md / AGENTS.md / README / LICENSE / VERSION / rendered HTML / rendered PDF are OUT of scope. Agents **MUST** read the ****Revision:**** N literal from the document head when reporting document state, **NEVER** infer freshness from filename mtime or commit log — both can lie. Agents writing to an IN-scope doc **MUST** invoke scripts/doc_revision_bump.sh <file> as the final step before staging, **OR** rely on the pre-commit hook. Agents **MUST NOT** manually edit the revision number — only the bump script is authoritative.

Pre-build gates CM-DOC-REVISION-HEADER-PRESENT + CM-COVENANT-114-44-PROPAGATION + paired mutations. Composes with §12.10 (CONTINUATION.md header reuse), §11.4.12, §11.4.22, §11.4.23, §11.4.18. No escape hatch. Classification: universal (§11.4.17). See Constitution §11.4.44 for the full mandate.

Integration-status-doc maintenance (§11.4.45, User mandate 2026-05-18)

Every non-trivial domain integration **MUST** have a docs/<domain>/<integration>/Status.md carrying the §11.4.44 revision header + auto-synced HTML+PDF + auto-colored per §11.4.23 + sync wrapper + captured-evidence table (every claim cites the test log / recording / sink-probe path) + closed status vocabulary (PASS / FAIL / SKIP / PENDING_FORENSICS / OPERATOR-BLOCKED) + operator-blocked items at the top + reference from CONTINUATION.md §3 when non-terminal. §11.4.45 is the generic form of §12.10 applied to every integration domain.

Pre-build gates CM-COVENANT-114-45-PROPAGATION + CM-AF-INTEGRATION-STATUS-DOCS + 4 paired mutations. Composes with §11.4.5 / §11.4.12 / §11.4.13 / §11.4.15 / §11.4.22 / §11.4.23 / §11.4.44 / §12.10. No escape hatch. Classification: universal (§11.4.17). See Constitution §11.4.45 for the full mandate.

Validate-recent-work-before-post-flash-tests (§11.4.46, User mandate 2026-05-18)

Every post-flash run **MUST** first run a recent-work validation pass (targeted tests for Issues.md In-progress / Ready-for-testing / Reopened + Fixed.md last-7-days + CONTINUATION.md §3 active work). Full test_all_fixes.sh runs **ONLY** after 100% green.

Helper: scripts/testing/recent_work_validate.sh --device <serial> writes /data/local/tmp/.recent_work_validated; full suite refuses without it. Each recent fix needs paired §11.4.43 RED-then-GREEN — naked GREEN is a bluff. Composes with §11.4.4 + §11.4.7 + §11.4.40 + §11.4.42 + §11.4.43 + §11.4.44 + §12.10. Pre-build gates CM-COVENANT-114-46-PROPAGATION + CM-AF-RECENT-WORK-VALIDATION-GATE + CM-AF-VALIDATION-ARTIFACT-FILE + paired mutations. Classification: universal (§11.4.17). No escape hatch. See Constitution §11.4.46 for the full mandate.

Firestore data review (§11.4.47, User mandate 2026-05-18)

Before every “bigger working round” (pre-build / pre-flash / pre-tag blocking; daily / post-deployment burn-in non-blocking) the operator/loop MUST execute `scripts/firebase/review_round.sh`. The pass queries Crashlytics (fatals + non-fatals + ANRs) + Analytics + Performance, classifies findings by severity, deduplicates to existing Issues.md entries via a three-tier algorithm (exact Firestore Issue-ID match → stacktrace-similarity cluster hash → operator merge review), and drafts new Issue entries for unrecognised findings with full Firestore Console URL refs + §11.4.4(a) systematic-debugging output.

Five mandatory elements: 5-trigger cadence, 3-source query (all three sources), Issues.md output with Firestore metadata (Issue IDs + URL + Cluster Hash / KPI / Funnel), 3-tier dedup, and comprehensive root-cause analysis per Issue.

Pre-build gates CM-COVENANT-114-47-PROPAGATION + CM-AF-FIREBASE-REVIEW-CADENCE + CM-AF-FIREBASE-ISSUE-XREF + 3 paired mutations. Composes with §11.4.4 / §11.4.4(a) / §11.4.6 / §11.4.7 / §11.4.10 / §11.4.12 / §11.4.14 / §11.4.15 / §11.4.16 / §11.4.34 / §11.4.42 / §11.4.43 / §11.4.44 / §11.4.45 / §11.4.46. No escape hatch — no `--skip-firebase-review`. Classification: universal (§11.4.17). See Constitution §11.4.47 for the full mandate.

UI-driven video testing (§11.4.48, User mandate 2026-05-18)

Every test that asserts video playback on a secondary display MUST traverse the user-equivalent UI path (launcher icon → app home → content list → tile tap → playback → in-app pause/resume → back button stop). NOT Intent/Broadcast shortcuts (`am start -a VIEW, cmd media_session play`). `Realuiautomator` dump element resolution + input `tap X Y`. Coverage: every video-capable app in `PRODUCT_PACKAGES` + every stream type (progressive HTTP / HLS / DASH / RTMP / file-local / DRM) + every codec (H.264/265/VP9/AV1/MPEG-2/MP4V video; AC-3/E-AC-3/TrueHD/DTS/DTS-HD/MLP/Opus/AAC audio). Secondary display verified via `ffprobe-on-captured-mp4` + VOM `activeDecoder` state. Arvus codec-state cross-check per §11.4.13 + §CG screenshot.

Per §11.4.4 four-layer: pre-build gate CM-AF-UI-DRIVEN-VIDEO- COVERAGE + propagation gate CM-COVENANT-114-48-PROPAGATION + on-device test framework at `device/rockchip/rk3588/tests/ui_driven/` (Layer 1 helper + Layer 2 per-app drivers + Layer 3 scenarios) + Layer 4 orchestrator `scripts/testing/run_ui_driven_video_suite.sh` + paired meta-test mutations. No escape hatch — no `--use-intent-shortcut`, `--skip-ui-traverse`, `--legacy-intent-mode` flag. Classification: universal (§11.4.17). See Constitution §11.4.48 for the full mandate.

Dual-approach testing (§11.4.49, User mandate 2026-05-18)

Every feature test exercising a user-visible behaviour MUST ship in TWO variants: a UI-driven variant (uiautomator-based, §11.4.48 surfaces A–E) AND an Intent/Broadcast-driven variant (`am start --es / am broadcast`-based). Either alone is a §11.4 PASS-bluff for the OPPOSITE half of the stack — UI catches app-side bugs; Intent catches framework/system-server bugs (MediaCodec.configure hook, IVideoOutputManager bind timing, broadcast permission gating).

Shared assertion base tests/lib/dual_approach_test_base.sh — both variants share dat_init / dat_start_capture / dat_assert_codec_state / dat_assert_video_frames / dat_assert_audio_channels / dat_arvus_dashboard_capture / dat_cleanup / dat_report_finding. Both variants gather identical evidence into mirror dirs qa-results/dual_approach/ <F>/<run-ts>/{ui,intent}/ so orchestrator diffs results and pinpoints which half of the stack contains a bug. Status mismatch between variants is itself a finding.

Kinopoisk 5.1 EAC3 is the canonical first implementation. Both variants are RED per §11.4.43 until the §CN decoder pipeline fix lands. Pre-build gates: CM-COVENANT-114-49-PROPAGATION + CM-AF-DUAL-APPROACH-COVERAGE + CM-AF-KINOPOISK-5-1-DUAL- COVERAGE. Three paired meta-test mutations. No escape hatch — no --ui-only / --intent-only / --skip-dual flag. Classification: universal (§11.4.17). See Constitution §11.4.49 for the full mandate.

Deterministic consistency (§11.4.50, User mandate 2026-05-18)

Every test that PASSES MUST have been executed N times (default N=3 normal tests, N=10 cycle-validation suites) against the same firmware MD5 + same device + same topology and produced IDENTICAL PASS in every iteration. A divergent N-iter run is auto-FAIL — there is no “first PASSED therefore X was a flake” path. §11.4.7’s intermittent / transient / flake / flaky vocabulary is enforced MECHANICALLY by this mandate, not merely textually.

Coverage: every public API path (Activity / Service / Broadcast receiver / ContentProvider URI / IPC interface / JNI entry / sysprop write / sysfs node / init.rc trigger) MUST have ≥1 dedicated test that drives it. Untested paths surface in a feature-coverage-matrix audit; the threshold ratchets 70 → 85 → 95 → 99 over phases.

Reliability mechanism: ab_run_n_times helper in the project anti-bluff library loops, captures evidence-hash per iter, asserts all N hashes + exit codes identical. NO operator-facing escape converts divergence to PASS.

Pre-build gates: CM-COVENANT-114-50-PROPAGATION + CM-AF-RELIABILITY-CHECK-WIRED + CM-AF-FEATURE-COVERAGE-MATRIX. Three paired meta-test mutations. No escape hatch — no --allow-flake, --first-pass-suffices, --skip-n-iter, --skip-coverage-audit flag. Classification: universal (§11.4.17). See Constitution §11.4.50 for the full mandate.

Live-ADB-First maximization (§11.4.51, User mandate 2026-05-18)

Every fix MUST be classified by rebuild-requirement before commit using the project’s per-file-class decision matrix. If LIVE_ADB_TESTABLE (on-device test scripts, host scripts, atmosphere-.sh boot scripts, persist. properties, markdown docs, test fixture assets, HelixQA YAML banks), the operator MUST first apply the fix to the running device via adb push / setprop / pm install -r / mount -o remount, rw, run the §11.4.43 RED test live, capture PASS, THEN commit + rebuild + reflash. Commit footer: LIVE_ADB_VALIDATED: yes. If REQUIRES_REBUILD (kernel, framework Java/AIDL, native C++ in APEX, sepolicy, init.rc, ro.* properties, XML overlays, codec XML in APEX, Android.bp/.mk), the operator proceeds directly to source-side + rebuild. Commit footer: REQUIRES_REBUILD: <reason>. Mixed batches use partial.

§11.4.51 REFINES §11.4.43 step 2 with mechanical enforcement. Helper: `classify_fix_rebuild_requirement.sh` walks `git diff --name-only`, looks up each file against the matrix, emits per- file classification + recommended commit-message footer. Unmatched paths classify as `REQUIRES_REBUILD: unmatched-path` (safe default per §11.4.6).

Pre-build gates: `CM-COVENANT-114-51-PROPAGATION` + `CM-AF-CLASSIFY-FIX-HELPER-EXISTS` + `CM-AF-LIVE-ADB-FIRST-COMMIT-MARKER`. Three paired meta-test mutations. No escape hatch — no `--skip-classify` / `--assume-rebuild` / `--no-footer-required` flag.

Classification: universal (§11.4.17). See Constitution §11.4.51 for the full mandate.

Autonomous-Validation (§11.4.52, User mandate 2026-05-18)

Forensic anchor — verbatim user mandate:

“Make sure we have full automation tests which will do all this work in full automation! IMPORTANT: Make sure that all existing tests and Challenges do work in anti-bluff manner — they MUST confirm that all tested codebase really works as expected! execution of tests and Challenges MUST guarantee the quality, the completion and full usability by end users of the product!”

Every user-facing feature MUST have at least one autonomous validation path: end-to-end via `adb shell` + scripted automation, captured runtime evidence per §11.4.5, PASS/FAIL verdict WITHOUT human presence. Operator-attended tests are SUPPLEMENTARY, never PRIMARY. A feature whose ONLY path is operator-attended is a §11.4.52 violation: the path does not scale to CI, does not run on every commit, does not survive operator unavailability, and produces the exact “tests pass but feature doesn’t work for users” failure mode §11.4 forbids.

Acceptable autonomous paths: instrumentation APK (SDK-API exercises + JSON result file), headless intent dispatch + state poll (`am start --es + dumpsys//proc/<pid>/maps` polling), ADB-driven `uiautomator` (ONLY if hierarchy has ≥ 1 clickable node — empty hierarchy demands fallback to APK/intent), network-side sink probe (§11.4.13), HelixQA autonomous QA session (§11.4.27).

Coverage ledger (§11.4.25) classifies each feature as `AUTONOMOUS_VERIFIED` / `AUTONOMOUS_DESIGNED` / `OPERATOR_ATTENDED_ONLY` / `NOT_APPLICABLE`. `OPERATOR_ATTENDED_ONLY` blocks release until migrated; cite a tracked work item per §11.4.15 + §11.4.16.

Pre-build gates `CM-COVENANT-114-52-PROPAGATION` + `CM-AF-AUTONOMOUS-PATH-PER-FEATURE`. Paired mutations. No escape hatch — no `--allow-operator-attended-only` / `--skip-autonomous-path` / `--manual-validation-suffices` flag.

Classification: universal (§11.4.17). See Constitution §11.4.52 for the full mandate.

Fixed_Summary parity (§11.4.53, User mandate 2026-05-18)

Forensic anchor — verbatim user mandate (2026-05-18T17:55Z):

“Note: Just like for Issues we have Issues_Summary, for Fixed we MUST HAVE Fixed_Summary - like all other docs: ALWAYS in sync and up to date and ALWAYS exported into the PDF and HTML!”

docs/Fixed_Summary.md is the symmetric short-form summary of docs/Fixed.md. MUST be regenerated whenever Fixed.md changes. HTML + PDF exports MUST travel with the markdown. Stale exports are §11.4.53 violations regardless of whether the underlying .md is correct. Same discipline as §11.4.12 Issues_Summary applied to Fixed.md.

Generator: scripts/testing/generate_fixed_summary.sh (canonical). Auto-sync wrapper: scripts/testing/sync_issues_docs.sh regenerates BOTH Issues_Summary AND Fixed_Summary in one shot, exports HTML + PDF, colorizes per §11.4.23, re-renders PDFs. MUST be invoked after any edit to Fixed.md. No --issues-only flag exists.

Sort order: closure date DESC (most-recent-Fixed first), §-letter / Fix-# secondary. Documented at top of generated file.

Composes with §11.4.12 (Issues_Summary sibling), §11.4.19 (atomic migration trigger), §11.4.23 (colorizer), §11.4.33 (type-aware closure terminal values), §11.4.44 (revision header), §12.10 (CONTINUATION.md resumption).

Pre-build gates CM-FIXED-SUMMARY-SYNC + CM-COVENANT-114-53-PROPAGATION. Paired mutations. No escape hatch — no --skip-fixed-summary-sync, --issues-only, --summary-not-applicable flag.

Classification: universal (§11.4.17). See Constitution §11.4.53 for the full mandate.

ATM-NNN ticket identifier (§11.4.54, User mandate 2026-05-19)

Forensic anchor — verbatim user mandate (2026-05-19):

“Every workable item in Issues / Issues_Summary / Fixed / Fixed_Summary MUST carry a stable, unique, auto-incremental ATM-NNN ticket identifier.”

Every workable item heading in docs/Issues.md AND docs/Fixed.md MUST carry [ATM-NNN] (form ## §X.Y. [ATM-NNN] <title>, zero-padded ≥3 digits). Allocated by scripts/testing/assign_atm_ticket_ids.sh and persisted to scripts/testing/.atm_ticket_state.jsonl (jsonl). Identifiers are monotonic, never renumbered, never reused, no gaps. Issues_Summary.md and Fixed_Summary.md MUST carry an ATM ID column as the leftmost data column.

Composes with §11.4.15 + §11.4.16 + §11.4.19 + §11.4.33 + §11.4.12 + §11.4.53 + §11.4.55 + §11.4.57.

Pre-build gates CM-ATM-TICKET-IDS-COMPLETE + CM-ATM-TICKET-IDS-UNIQUE + CM-ATM-TICKET-IDS-MONOTONIC + CM-COVENANT-114-54-PROPAGATION. Paired mutations. No escape hatch.

Classification: universal (§11.4.17). See Constitution §11.4.54 for the full mandate.

Reopens-history tracking + per-item Reopens.md (§11.4.55, User mandate 2026-05-19)

Forensic anchor — verbatim user mandate (2026-05-19):

“Add a Reopens-count column. For any item whose reopens-count > 0, create docs/issues/ATM-NNN/Reopens.md (+ HTML + PDF) with comprehensive reopen history + Fixed cycles.”

Every workable item with `reopens_count > 0` MUST have docs/issues/<ATM-NNN>/Reopens.md (+ HTML + PDF). Required content: §11.4.44 revision header, item identification, cycle counters, chronological timeline (By AI/User per §11.4.34, On ISO date, Event, Reason, Evidence, Outcome), reasoning chain for each closure, most-recent state-change pointer.

Issues_Summary.md and Fixed_Summary.md MUST carry a Reopens column; count > 0 hyperlinks to per-item Reopens.md.

Composes with §11.4.34 (per-event capture), §11.4.54 (ATM-NNN path), §11.4.44 (revision header), §11.4.45 (Status.md analogue), §11.4.53 (Fixed_Summary parity — Reopens column symmetric).

Pre-build gates CM-REOPENS-DOC-EXISTS-WHEN-COUNT-GT-ZERO + CM-REOPENS-DOC-REVISION-HEADER + CM-REOPENS-COL-IN-SUMMARIES + CM-COVENANT-114-55-PROPAGATION. Paired mutations. No escape hatch.

Classification: universal (§11.4.17). See Constitution §11.4.55 for the full mandate.

Status_Summary parity + two-audience format (§11.4.56, User mandate 2026-05-19)

Forensic anchor — verbatim user mandate (2026-05-19):

“Every Status.md doc gets a Status_Summary parity companion ALWAYS in sync, exported to HTML + PDF. Two-page format: page 1 = non-developer audience (team-specific), page 2 = software engineer summary.”

For every docs/<domain>/<integration>/Status.md a companion Status_Summary.md MUST exist with: §11.4.44 revision header; **Page 1 — For the** (audience-specific — audio team for docs/dolby/*, video team for docs/video/*, etc.) plain- language summary, What works, What’s broken or pending, Operator actions — NO §-letter jargon, NO captured-evidence file paths; **Page 2 — For software engineers** — §-letter refs, gate names, commit hashes, captured-evidence paths, ATM-NNN cross-refs. HTML + PDF exports travel with the markdown.

Generator: scripts/testing/generate_status_summary.sh <Status.md path>. Invoked from sync_integration_status.sh (§11.4.45 wrapper) on every Status.md update.

Composes with §11.4.45 (Status.md — Status_Summary COMPLEMENTS), §11.4.12 + §11.4.53 (parity discipline), §11.4.44 (revision header), §11.4.23 (colorizer), §12.10 (CONTINUATION.md).

Pre-build gates CM-STATUS-SUMMARY-EXISTS-FOR-EVERY-STATUS + CM-STATUS-SUMMARY-TWO-AUDIENCE + CM-STATUS-SUMMARY-REVISION-HEADER + CM-COVENANT-114-56-PROPAGATION. Paired mutations. No escape hatch.

Classification: universal (§11.4.17). See Constitution §11.4.56 for the full mandate.

README.md doc-link section + revision metadata (§11.4.57, User mandate 2026-05-19)

Forensic anchor — verbatim user mandate (2026-05-19):

“Add a doc-link section to README.md — links to Issues + Issues_Summary + Fixed + Fixed_Summary + CONTINUATION + ALL Status docs + their exports. Each link shows revision + last-modified.”

Every project's top-level README.md MUST contain a section titled Tracked-Items + Status Documents (heading MUST contain literal Tracked-Items). The section is a markdown table with columns: Document, Last modified, Revision, Markdown, HTML, PDF. Links to: Issues.md + Issues_Summary.md, Fixed.md + Fixed_Summary.md, CONTINUATION.md, every docs/**/Status.md + its Status_Summary.md pair. Status docs auto-discovered by the generator via `find docs -name 'Status.md'`.

Generator: `scripts/testing/update_readme_doc_links.sh` extracts each doc's revision + last-modified from its §11.4.44 header, renders the table, replaces the section between explicit markers (`<!-- doc-link-section:begin -->` / `<!-- doc-link-section:end -->`). Invoked from `sync_issues_docs.sh` AND `sync_integration_status.sh`.

Composes with §11.4.12 + §11.4.19 + §11.4.53 (Issues / Fixed / summaries), §11.4.44 (revision-header data source), §11.4.45 + §11.4.56 (Status pairs), §12.10 (CONTINUATION.md).

Pre-build gates CM-README-DOC-LINK-SECTION-PRESENT + CM-README-DOC-LINK-ROWS-COMPLETE + CM-README-DOC-LINK-FRESHNESS + CM-COVENANT-114-57-PROPAGATION. Paired mutations strip markers, remove rows, backdate Last modified, strip anchor literal. No escape hatch.

Classification: universal (§11.4.17). See Constitution §11.4.57 for the full mandate.

Parallel-development methodology (§11.4.58, User mandate 2026-05-19)

Forensic anchor — verbatim user mandate (2026-05-19T~05:00Z MSK):

“We MUST CREATE adjusted improved version of working methodology where multiple workable items could be done in parallel, all come into the central (main) branch as they are done, then we at particular moment rebuild and reflash the System and in background full testing with validation and verification is done. Each parallel work on one workable (or more) item(s) must use parallel agents as much as possible! ... Writing in depth tests (all supported types of the tests) with Challenges and full HelixQA use is

MANDATORY! Every test we execute besides executed with success MUST RESULT in proof that actual functionality being tested REALLY DOES WORK with NO BLUFF of any kind!”

Project work proceeds through the **Parallel Work Unit (PWU) pipeline**. Each PWU is a self-contained workable item with: ATM-NNN identifier per §11.4.54, Issues.md entry per §11.4.15+§11.4.16, file- scope manifest, §11.4.43 RED test, source patch, pre-build gate, post-flash test, paired meta-test mutation per §1.1, HelixQA Challenge bank entry, captured-evidence directory per §11.4.5+§11.4.52.

5-stage pipeline: Stage 1 DEVELOP (parallel PWU agents in isolated worktrees) → Stage 2 MERGE (serial via `commit_all.sh` flock + §11.4.41 4-step merge-first) → Stage 3 REBUILD+FLASH (parallel where hardware allows) → Stage 4 VALIDATE (parallel D3+D4+meta-test+coverage) → Stage 5 SWEEP (parallel HelixQA + Fixed.md migration + README refresh). Stage 1 of round N+1 overlaps with Stages 4-5 of round N — throughput multiplier.

Synchronization: 4-layer lock hierarchy (parent flock / per- submodule git / contention-path advisory locks / per-PWU worktree). Disjoint-scope PWUs fully parallel.

Anti-bluff merge-time enforcement (mandatory, all four): RED-test captured (§11.4.43), paired meta-test mutation (§1.1), 3-iter deterministic-consistency (§11.4.50), captured-evidence per §11.4.5. Metadata-only / configuration-only / absence-of-error PASS REJECTED. HelixQA Challenge MANDATORY for every user-visible PWU.

Pre-build gates CM-PWU-LOCK-HIERARCHY + CM-PWU-ANTI-BLUFF-COVERAGE + CM-PWU-MERGE-QUEUE-DISCIPLINE + CM-PWU-PARALLEL-AGENT-LIMIT + CM-COVENANT-114-58-PROPAGATION. Paired mutations cover each gate. No escape hatch.

Classification: universal (§11.4.17). See Constitution §11.4.58 for the full mandate. Project-specific implementation reference in consumer-side docs/guides/ PARALLEL_DEVELOPMENT_METHODODOLOGY.md.

README always-sync (§11.4.59, User mandate 2026-05-19)

README.md is a §11.4.12-class always-sync document. HTML + PDF exports refresh on every update via `scripts/testing/sync_readme_export.sh` (pandoc + weasyprint); auto-invoked by `sync_issues_docs.sh` so a single doc-sync run refreshes Issues / Issues_Summary / Fixed / Fixed_Summary / CONTINUATION / README (md + html + pdf). README carries a §11.4.44 revision header and a Documentation Map section linking to every Status.md + Status_Summary.md + spec + plan + guide + script-companion doc + changelog + the constitution submodule, plus per-audience navigation (end user / developer / QA / agent). Pre-build gate CM-README-EXPORT-SYNC enforces mtime parity (README.html + README.pdf ≥ README.md). Paired meta-test mutation backdates HTML+PDF → gate FAILs. No escape hatch — no `--skip-readme-sync`, `--no-readme-export`, `--readme-stale-OK` flag. Composes with §11.4.12 + §11.4.18 + §11.4.44 + §11.4.45 + §11.4.53 + §11.4.56 + §11.4.57 + §12.10.

Classification: universal (§11.4.17). See Constitution §11.4.59 for the full mandate.

Documentation always-sync composite covenant (§11.4.60, User mandate 2026-05-19)

Eight doc classes — Issues, Issues_Summary, Fixed, Fixed_Summary, CONTINUATION, README, every Status.md (domain-scoped), every Status_Summary.md — MUST be in sync across .md + .html + .pdf. Composite pre-build gate CM-DOCS-COMPOSITE-SYNC walks all 8 classes (Status fleet via recursive docs/** find), FAILs the build if ANY .html or .pdf mtime is older than its .md. Per-class anchors §11.4.12 / §11.4.44 / §11.4.45 / §11.4.53 / §11.4.56 / §11.4.57 / §11.4.59 / §12.10 govern individually; the composite catches what they miss when one is silently bypassed. Paired mutation backdates docs/Issues.html → gate FAILs. No escape hatch.

Classification: universal (§11.4.17). See Constitution §11.4.60 for the full mandate.

Credentials-handling (§11.4.10)

Forensic anchor — verbatim user mandate (2026-05-12):

| “Credentials or any secret and sensitive data MUST NOT leak!”

.env patterns git-ignored project-wide. Runtime-load-only — tests load credentials at execution time from operator-populated files under scripts/testing/secrets/ (or per-project equivalent). Test scripts MUST NEVER print or log credentials.

Host-session safety (§12)

Project procedures MUST NOT use more than **60%** of total system RAM. Heavy work wrapped in bounded execution scopes so the kernel OOM-kills only the scope — user@<uid>.service stays alive. Hibernating / suspending / logging out the user is FORBIDDEN (directly OR indirectly).

Continuation document (§12.10)

docs/CONTINUATION.md MUST exist at the project root and reflect the live state of the work. Every non-trivial state change updates it in the same commit. Any agent must be able to resume work exactly where the previous session left off by reading this single file.

Data safety (§9)

Hardlinked backup before any destructive operation. Post-op gate MUST pass. Force-push requires explicit per-session human authorization. Every history rewrite gets a docs/change logs/ audit entry.

CLI workflow expectations

1. **Read before write.** Use file-reading tools to inspect the current state before editing. Edits in the dark are how bugs sneak in.

2. **Use the project's commit wrapper.** Direct `git add / git commit` is forbidden unless explicitly carved out.
3. **Commit messages cite forensic evidence** for any non-trivial fix (logs, captures, /sys readings, dropbox entries, strace). Speculative narratives are §11.4.6 violations.
4. **Update CONTINUATION.md** in the same commit as the work (§12.10).
5. **Run pre-build / post-build / runtime gates** before declaring a change complete. NEVER report “done” without running the relevant gate.
6. **Cite external research sources** for non-trivial fixes (§11.4.8). Either a citation OR the literal “NO external solution found — original work”.
7. **Refuse to bluff.** If a step requires information you don't have, say so and STOP. Don't guess.

Hierarchy of project authority

When a project files declares conflicting rules with this base AGENTS.md, resolve in this order:

1. **Project's docs/guides/<PROJECT>_CONSTITUTION.md** — most specific, project-tailored.
2. **constitution/Constitution.md** (this submodule) — universal.
3. **Project's root CLAUDE.md / AGENTS.md** — agent-facing summaries; lossy compared to the Constitution.
4. **This constitution/CLAUDE.md / AGENTS.md** — universal agent-facing summary; lossy compared to the Constitution.

If a project's own files contradict the Constitution, that is itself a Constitution violation and **MUST** be reported. Do NOT silently follow the contradicting project file.

When stuck

- Read constitution/Constitution.md.
- Cross-reference the project's CLAUDE.md / AGENTS.md.
- If still unclear, ask the operator. Do NOT guess. Do NOT bluff.

§11.4.65 — Universal Markdown export mandate (User mandate, 2026-05-19)

Every Markdown document inside the project that is NOT part of an application or service's source-code tree **MUST** have synchronized `.html` and `.pdf` siblings. Includes: project-root `*.md`, docs/**/*`.md`, scripts/**/*`.md` (doc-format companion docs), owned-submodule top-level README.md / CLAUDE.md / AGENTS.md / CHANGELOG.md and their docs/**/*`.md`, constitution/**/*`.md`, owned HelixQA submodules' equivalents. Excludes: external/**, prebuilts/**, packages/modules/**, kernel-5.10/**, out/**, build/**, application/service source-code trees, and third-party submodules NOT in the owned set. Every edit triggers regeneration via `scripts/testing/sync_all_markdown_exports.sh` (pandoc HTML + weasyprint PDF, timeout 60 per file, capped at 500 candidates). HTML + PDF mtime **MUST** be $\geq$ source `.md` mtime at all times.

Pre-build gates CM-UNIVERSAL-MARKDOWN-EXPORT-SYNC + CM-COVENANT-114-65-PROPAGATION. Paired meta-test mutations. Composes with §11.4.12 / §11.4.18 / §11.4.23 / §11.4.44 / §11.4.45 / §11.4.53 / §11.4.59 / §11.4.60 / §11.4.63 / §11.4.64. No escape hatch — no --skip-md-exports, --no-pdf-only, --md-export-not-applicable flag.

Canonical authority: constitution submodule Constitution.md §11.4.65.

Non-compliance is a release blocker regardless of context.

§11.4.66 — Blocker-resolution interactive-clarification mandate (User mandate, 2026-05-19)

When any task is blocked (operator decision, hardware access, external authorization, ambiguous scope), the agent MUST: (1) research what's doable from the agent side without operator input; (2) calculate minimum-viable operator input; (3) construct 2–4 mutually-exclusive options with one marked “Recommended” and each stating what the agent does after that answer; (4) present via the platform's interactive question mechanism (AskUserQuestion on Claude Code) — NEVER free-text “what would you like?” for closed- set decisions; (5) after the answer, resume work without follow-up round-trips. Composes with §11.4.6 / §11.4.7 / §11.4.40 / §11.4.41 / §11.4.42 / §11.4.52. No silent waiting; no bulk-text questions when interactive options would do.

Pre-build gate CM-COVENANT-114-66-PROPAGATION enforces the anchor literal across the 42-file consumer fleet. Paired meta- test mutation strips the literal → gate FAILs. No escape hatch — no --skip-ask, --silent-wait, --free-form-only flag.

Canonical authority: constitution submodule Constitution.md §11.4.66.

Non-compliance is a release blocker regardless of context.

§11.4.67 — Shell-script target-shell-parseability mandate (User mandate, 2026-05-19)

Forensic anchor — direct user mandate (verbatim, 2026-05-19): “any issue we spot must be fixed, bash scripts as well if they are broken!” + “Make sure that this is mandatory rule!”

Every shell script that may be invoked under a target shell other than the one in its shebang MUST parse cleanly under that target shell. Forensic incident: device/rockchip/rk3588/tests/test_all_fixes.sh:114 used bash-only `exec >>(tee -a "$f") 2>&1` on a sh script.sh callsite — Android mksh parses the whole script BEFORE executing, so the runtime `[ -n "${BASH_VERSION:-}" ]` guard could not save it. Fixed by wrapping in `eval 'exec >>(tee ...) 2>&1'` so the parser sees only a string.

Closed-set scope: every tracked .sh under device/rockchip/rk3588/tests/, scripts/, scripts/testing/ (and equivalent paths in owned submodules). OUT of scope: external/, prebuilts/, packages/modules/, kernel-5.10/, out/, build/, scripts/legacy/. Mandatory invariants: (1) every in-scope script parses under `sh -n`; (2) bash-only constructs (`>(…)`, `<(…)`, `[[ ]]`, `<<<`, arrays, `${var^^}`, etc.) MUST be wrapped in `eval` OR guarded by bash-only loading; (3)

shebangs honest — `#!/bin/bash` only if bash actually expected; (4) fix at source per §11.4.1, never at callsites. Composes with §11.4.1 / §11.4.4 / §11.4.6 / §11.4.50 / §11.4.51.

Pre-build gate `CM-SCRIPT-TARGET-SHELL-PARSEABLE` runs `sh -n` on every in-scope script. Propagation gate `CM-COVENANT-114-67-PROPAGATION` enforces the anchor literal across the 44-file consumer fleet. Paired mutations: inject bash-only outside `eval` → parse gate FAILs; strip `11.4.67` literal → propagation gate FAILs. No escape hatch — no `--skip-parseability-check`, `--bash-only-script`, `--runtime-guard-suffices` flag.

Canonical authority: constitution submodule [Constitution.md](#) §11.4.67.

Non-compliance is a release blocker regardless of context.

§11.4.68 — Positive sink-side / downstream evidence mandate (User mandate, 2026-05-20)

A test asserting audio or video routing PASS MUST capture and verify positive sink-side or downstream evidence — never config-only, never metadata-only, never PCM-open-state-only. Closed enumeration of acceptable evidence: (1) sink-side codec-state showing non-empty codec matching the expected regex during playback; (2) PCM `hw_ptr` delta strictly positive across the playback window; (3) ALSA ELD demonstrating negotiated channel count + format; (4) `ffprobe` non-zero frames matching expected codec for video; (5) recording-analyzer matched event per §11.4.2 / §11.4.5 timeline; (6) `tinycap` RMS amplitude above floor.

Failure modes specifically forbidden: silent SKIP on sink unreachable, empty codec-state treated as evidence, policy-side device field used as proof PCM opened, config-XML-only PASS. All produce the §11.4 PASS-bluff failure mode anchored at D3 forensic 2026-05-20.

Mandatory protections: sink-side libraries expose `*_require_reachable` returning exit code 2 (OPERATOR-BLOCKED); harness propagates 2; audio/video-routing tests capture ≥1 positive evidence per enumeration; anti-stickiness post-stop re-probe. No escape hatch.

Composes with §11.4.2 / §11.4.5 / §11.4.13 / §11.4.14 / §11.4.46 / §11.4.49 / §11.4.50 / §11.4.52. Pre-build gates `CM-COVENANT-114-68-PROPAGATION` + `CM-POSITIVE-SINK-EVIDENCE-PER-AUDIO-TEST` + paired meta-test mutations per §1.1.

Canonical authority: constitution submodule [Constitution.md](#) §11.4.68.

Non-compliance is a release blocker regardless of context.

§11.4.69 — Universal sink-side positive-evidence taxonomy + mechanical enforcement (User mandate, 2026-05-20)

Forensic anchor — direct user mandate (verbatim, 2026-05-20):

“THIS MUST HAPPEN NEVER AGAIN!!! We MUST HAVE this all working! Not just for audio but for every single piece of the System!!! Proper full automation when executed with success MUST MEAN that manual testing will be as much positive at least regarding the success results! ... Solution MUST BE universal, generic that solves working flows for all

System components and for all future and all existing projects! ... Everything we do MUST BE validated and verified with rock-solid proofs and anti-bluff policy enforcement and fulfillment!”

Universal generalisation of §11.4.68 (audio-specific) across every user-visible feature class. Closes the PASS-bluff pattern where tests reported green while end users hit broken features (2026-05-19→20 D3 audio “82/84 PASS” + empty Arvus Codec-In-Use).

The mandate. Every user-visible feature MUST map to one entry in the closed-set §11.4.69 sink-side evidence taxonomy (audio_output, audio_input, video_display, network_throughput, network_connectivity, bluetooth_a2dp, bluetooth_pair, touch_input, sensor, gpu_render, storage_read, storage_write, mediacodec_decode, mediacodec_encode, miracast, cast, boot_service, package_install, permission_grant, wifi_link, wifi_throughput, ethernet_link, display_topology, drm_playback, subtitle_render — open to additions). Every PASS for a feature in the taxonomy MUST cite a captured-evidence artefact path matching the required evidence shape.

Helper contracts (additive during grace; mandatory after 2026-06-19):

- `ab_pass_with_evidence <description> <evidence_path>` — the new canonical PASS helper. Verifies path exists AND non-empty; emits PASS: `<description> [evidence: <path>]`.
- `ab_skip_with_reason <description> <closed-set-reason>` — reasons: `geo_restricted`, `operator_attended`, `hardware_not_present`, `topology_unsupported`, `network_unreachable_external`, `feature_disabled_by_config`. Forbids `network_unreachable_external` for any taxonomy feature with a sink-side probe.
- Bare `ab_pass` deprecated — WARN pre-grace, FAIL post-grace (2026-06-19).

Mechanical enforcement. Three pre-build gates + three paired §1.1 meta-test mutations:

- `CM-SINK-EVIDENCE-PER-FEATURE` — walks tests for # §11.4.69 `FEATURE: <class>` annotation + verifies taxonomy probe + `ab_pass_with_evidence` use.
- `CM-NO-FAIL-OPEN-SKIP` — audits sink-side probe helpers; FAILs if any code path converts empty/unreachable response to PASS-counting SKIP for a feature class with a sink-side probe.
- `CM-AB-PASS-WITH-EVIDENCE-EVERYWHERE` — pre-grace WARN, post-grace FAIL on bare `ab_pass` calls.

Composes with §11.4.1 (FAIL-bluffs forbidden), §11.4.2 (recorded-evidence), §11.4.5 (audio + video 5-layer quality), §11.4.6 (no-guessing), §11.4.13 (sink-side captured-evidence), §11.4.27 (no-fakes-beyond-unit), §11.4.50 (deterministic consistency), §11.4.52 (autonomous-validation), §11.4.68 (audio-specific sink-side — §11.4.69 is the universal generalisation).

No escape hatch — no `--skip-evidence`, `--config-only-pass`, `--allow-fail-open-skip`, `--legacy-ab-pass-permitted` flag. The discipline exists because the 2026-05-20 forensic incident demonstrated the failure: tests reported audio-routing PASS while the user heard nothing and the Arvus Codec-In-Use field was empty.

Propagation gate CM-COVENANT-114-69-PROPAGATION enforces this anchor literal across the ~44-file consumer fleet. Paired mutation strips the literal → gate FAILs.

Canonical authority: constitution submodule Constitution.md §11.4.69.

Non-compliance is a release blocker regardless of context.

§11.4.70 — Subagent-driven execution is the default (User mandate, 2026-05-20)

Forensic anchor — direct user mandate (verbatim, 2026-05-20):

“Always do if possible Subagent-driven! Add this into our root (constitution Submodule) Constitution.md, CLAUDE.md and AGENTS.md. This should be the default choice ALWAYS!”

When executing implementation plans authored via `superpowers:writing-plans` (or any equivalent task-decomposed execution flow), the **default execution model is subagent-driven** per `superpowers:subagent-driven-development`. Inline execution via `superpowers:executing-plans` is permitted ONLY when (a) the task is trivial AND fits in a single sub-300-line edit, OR (b) the operator explicitly requests inline execution at brainstorm-handoff time.

Subagents bring an isolated context window per task (no conductor context bloat), a structurally separated review seam (conductor reviews subagent output, eliminating self-review blind spots), parallel-PWU compatibility (§11.4.58 — subagents ARE the parallel work units), and resumability across operator absence (subagents resume from on-disk plan + spec inputs).

Composes with §11.4.4 (four-layer coverage), §11.4.6 (no-guessing — subagent’s captured output IS the evidence), §11.4.42 (iteration discipline), §11.4.43 (TDD-fix), §11.4.50 (deterministic consistency), §11.4.51 (LIVE_ADB_FIRST), §11.4.58 (parallel- development PWU).

No escape hatch — `--inline-execution-required`, `--no-subagents`, `--monolithic-execution` are NOT permitted flags. Skipping subagent-driven for non-trivial work without recorded operator authorisation is itself a §11.4 PASS-bluff. Pre-build gate CM-COVENANT-114-70-SUBAGENT-DEFAULT-PROPAGATION enforces this anchor literal across the ~44-file consumer fleet. Paired meta-test mutation strips the literal → gate FAILs.

Canonical authority: constitution submodule Constitution.md §11.4.70.

Non-compliance is a release blocker regardless of context.

§11.4.71 — Pre-Push Fetch + Investigate + Integrate Mandate (User mandate, 2026-05-20)

Everyday-push variant of §11.4.41 (force-push merge-first). Before pushing to any upstream for any repository (main repo or submodule), the agent MUST follow the 5-step pre-push cycle: (1) `git fetch --all --prune --tags`, (2) `git pull --no-rebase <remote> <branch>` for each remote whose tip differs from local, (3) investigate the diff vs OUR previous HEAD by reading every foreign commit’s body (what changed, why, how it affects OUR project), (4) integrate mandatory

changes with full §11.4.4(b) four-layer anti-bluff test coverage producing REAL captured-evidence proofs, (5) THEN push to every configured remote in cascade order.

Composes with §11.4.41 (force-push case) + §11.4.26 / §11.4.32 / §11.4.37 / §11.4.40 / §11.4.42 / §11.4.43 / §11.4.4(b) / §11.4.5 / §11.4.6.

Applies to parent repo + constitution submodule + every owned submodule + every nested submodule + every HelixQA dependency. Audit-trail per push reconstructable from docs/change-logs/<tag>.md + per-repo git log evidence.

No escape hatch — no --skip-fetch, --no-investigate, --fast-push, --trust-upstream flag. Pre-build gate CM-COVENANT-114-71-PROPAGATION + paired mutation.

Canonical authority: constitution submodule Constitution.md §11.4.71.

Non-compliance is a release blocker regardless of context.

§11.4.72 — Audio Top-Priority Mandate (User mandate, 2026-05-20)

Direct user mandate (verbatim): “Make sure all fixes for audio are always top priority in main working stream!” The conductor (main working stream — Claude Code session, AI agent, or human operator) MUST treat audio fixes as the highest-priority class on the serial dispatch queue. Audio scope includes §EU HDMI rejection, §EM multichannel HDMI, §ET Arvus integration, §EV/§EW/§EX D3 audio defects, HiFi (Fix #74, #76), AC3 (Fix #106), ES8388 (Fix #103/§104/#105), multichannel LPCM (Fix #112), and every future audio-stack improvement.

Composes with §11.4.42 (iteration-discipline — audio sits at apex of priority order) + §11.4.58 (parallel-development PWU — background research subagents run concurrently but do NOT preempt audio on the main-stream serial dispatch queue).

Operative rule: any time the conductor faces a choice between dispatching an audio task vs a non-audio task on the SAME serial resource, the audio task wins. No escape hatch — no “but this non-audio task is faster” or “but this research is more interesting” override. Audio-stack regressions are user-perceptible and high-impact (D3 silent post-flash is a release blocker), while research and refactors can wait.

Canonical authority: constitution submodule Constitution.md §11.4.72.

Non-compliance is a process violation regardless of context.

§11.4.73 — Main-specification document versioning + revision discipline (User mandate, 2026-05-20)

Direct user mandate (verbatim): “Make sure everything we add now in previous and upcoming requests IS ALWAYS applied to the main specification — if we have one. Since all these are not major changes we could increase Specification version per change for secondary version instead of the primary. Primary version MUST BE increased for much bigger levels of changes! Document MUST BE updated ALWAYS to follow the versioning rules we are applying here + revision and other properties we have!”

Applies only when a project recognises a main specification document. Two-axis versioning: **primary** (V1/V2/V3/...) bumps for major rewrites (old primary archived); **secondary** (Revision) bumps for additive operator requirements (matches the §11.4.61 metadata-table Revision integer). Every operator-mandated requirement MUST land in the spec as part of the work that implements it. Cross-doc propagation copies MUST reference the active spec file, not a stale archived version. Composes with §11.4.44, §11.4.61, §11.4.59, §11.4.65.

Canonical authority: constitution submodule Constitution.md §11.4.73.

Non-compliance is a release blocker.

§11.4.74 — Submodule-catalogue-first discovery + extend-don't-reimplement (User mandate, 2026-05-20)

Direct user mandate (verbatim): “We MUST ALWAYS check which already developed features / functionalities do exist as a part of our comprehensive Submodules catalogue located in vasic-digital and HelixDevelopment organizations on GitHub and GitLab both! Project MUST BE aware of all its existence so we do not implement same things multiple times. For any missing features we MUST IMPLEMENT them properly and extend those Submodules further!”

Before scaffolding any new module / package / helper / utility, the agent MUST: (1) survey vasic-digital + HelixDevelopment orgs on GitHub + GitLab — the canonical inventory is submodules-catalogue.md (142 repos categorised); (2) reuse an existing Submodule when it covers $\geq 80\%$ of the functionality; (3) extend in-place via upstream PR when $80\%+$ matches but features are missing — never duplicate; (4) document the survey result in the relevant tracker row with Catalogue-Check: reuse|extend|no-match <org/repo>@<sha>.

Every Submodule in the catalogue is subject to the same development-cycle rules (§11.4 anti-bluff, §1.1 paired mutations, §11.4.10 credentials, §11.4.61 metadata + ToC, §11.4.65 universal export, §11.4.73 spec versioning, §2.1 multi-mirror push, §3 propagation order).

Canonical authority: constitution submodule Constitution.md §11.4.74.

Non-compliance is a process violation; severe cases (duplicate implementation landed without catalogue check) are release blockers.

§11.4.75 — Mechanical Enforcement Without Exception (User mandate, 2026-05-20)

Direct user mandate (verbatim): “Why do these violations still happen!? This is a serious problem! We cannot rely on stability nor consistency if we cannot respect our Constitution, mandatory rules and constraints! Is there a way to make this always respected, followed and applied without exception fully and unconditionally!? WE MUST HAVE THIS WORKING FLAWLESSLY!!! Do investigate the root causes of such problems! Once all problems are identified WE MUST apply proper mechanisms for this not to happen NEVER EVER AGAIN!”

The §11.4 covenant historically relied on agent + operator vigilance. Three forensic incidents in 2026-05-19→20 (two stalled remediation subagents + post-§11.4.66 propagation drift requiring catch-up commit f93b25a92eb) demonstrated that late-

binding enforcement at `pre_build_verification.sh` time fires hours-to-days AFTER the violator commit has already reached every remote. §11.4.75 closes the gap with FIVE independent mechanical enforcement layers — bypassing any single layer does not bypass the discipline:

1. **Local pre-commit git hook** — refuses staged `.md` lacking sibling `.html+.pdf` (and other staged-only invariants).
2. **commit_all.sh integration** — the canonical project commit script invokes the same checks + auto-runs `sync_all_markdown_exports.sh` to self-repair before commit (`_constitution_sibling_check` function).
3. **Local pre-push git hook** — re-runs siblings + propagation-gate subset on every commit in the push.
4. **post-commit auto-repair hook** — detects orphan `.md` in just-committed manifest, auto-generates siblings via `pandoc + weasyprint`, creates a chore(§11.4.75): `auto-export ... follow-up commit`. Idempotent + recursion-guarded.
5. **Local-only equivalent** (Phase 39.GF, User mandate 2026-05-20). Remote CI surfaces (GitHub Actions, GitLab pipelines, Jenkins, CircleCI, etc.) are DISABLED — the workflow file is preserved at `.github/workflows/constitution-compliance.yml.disabled-local-only` (NOT `.yml`; GitHub Actions ignores it). Layer 5 enforcement migrated to the LOCAL pre-build-verification + meta-test ritual the operator MUST execute before tagging per §11.4.40. Layers 1-4 remain authoritative; Layer 5 is the operator's local final gate. A future re-enable PWU may re-establish remote CI.

Helper contracts (mandatory): `scripts/install_git_hooks.sh` (idempotent installer wired into `scripts/setup.sh`), `scripts/git_hooks/{pre-commit,pre-push,post-commit,commit-msg}`, `_constitution_sibling_check` in `scripts/commit_all.sh`. The `commit-msg` hook enforces a Bypass-rationale: `<reason> footer` when `--no-verify` is detected (touch-file marker `.git/ATMO_LAST_BYPASS_ATTEMPT`); `docs/audit/bypass_events.md` accumulates the audit trail per §11.4 captured-evidence requirement.

Five pre-build gates with paired §1.1 meta-test mutations: CM-COVENANT-114-75-PROPAGATION (anchor literal across canonical files), CM-GIT-HOOKS-INSTALL-SCRIPT (installer present + executable), CM-GIT-HOOKS-SOURCE-DIR (4 hook bodies present + executable), CM-COMMIT-ALL-SIBLING-CHECK (`_constitution_sibling_check` in `commit_all.sh`), CM-CI-WORKFLOW-PRESENT (CI workflow + ≥ 3 required jobs).

Composes with §1.1 (paired meta-test mutations), §9 (data safety), §11.4 (end-user-quality covenant — Layer 5 CI makes the §11.4 promise mechanical), §11.4.41 (this anchor's renumber from §11.4.74 → §11.4.75 was itself driven by §11.4.41 merge-first discipline), §11.4.65 (universal Markdown export — Layer 1+3+4 are §11.4.65's mechanical seam), §11.4.66 (interactive clarification), §11.4.67 (target-shell-parseability — hooks parse under `bash AND mksh`), §11.4.71 (pre-push fetch + integrate — Layer 3 + 5 honour the merge-first pipeline), §11.4.72 (audio top-priority — hooks hold no lock themselves, so do not race against in-flight audio commits), §11.4.73 (main-spec versioning — when spec exists, hooks include it), §11.4.74 (submodule-catalogue-first — hooks can be added to the canonical catalogue for cross-project reuse).

No escape hatch — no `--skip-hooks`, `--bypass-enforcement`, `--allow-orphan-md`, `--ci-not-applicable`, `--mechanical-enforcement-not-needed` flag exists. The `--no-verify` route IS the deliberate audit-trail bypass; §11.4.75 makes the audit trail mechanical via the `commit-msg` footer requirement.

Canonical authority: constitution submodule Constitution.md §11.4.75.

Non-compliance is a release blocker regardless of context.

§11.4.76 — Containers-submodule mandate (User mandate, 2026-05-20)

Direct user mandate (verbatim): “For any work or requirements of running services or codebase inside the Containers (Docker / Podman / Qemu / Emulators, and so on) we **MUST USE / INCORPORATE** the Containers Submodule properly: <https://github.com/vasic-digital/containers> (`git@github.com:vasic-digital/containers.git`). Containers Submodule contains all means for us to Containerize our code and services! If any feature or Containing System is missing or not supported we **MUST EXTEND IT** properly like we do all of our projects! No bluff work is allowed of any kind!”

§-slot history note. Originally drafted as §11.4.75, renumbered to §11.4.76 per §11.4.71 fetch-before-push after the concurrent `0a70083` landing of §11.4.75 (Mechanical Enforcement).

For ANY containerized workload (Docker / Podman / Qemu / Kubernetes / container-backed emulators), the agent **MUST**: (1) install `vasic-digital/containers` (`digital.vasic.containers`) as a Git submodule when scaffolding the consuming project; (2) consume via `replace` directive during development + pinned commit SHAs in production; (3) boot infra on-demand via the Submodule’s `pkg/boot` + `pkg/compose` + `pkg/health` APIs so the operator is never required to start `podman machine / docker compose` up manually — the boot is part of the test entry point (**on-demand-infra invariant**); (4) extend the Submodule via upstream PR when a runtime / lifecycle primitive is missing — never reimplement in-project (per §11.4.74); (5) anti-bluff: integration tests claiming to exercise containerized components **MUST** actually boot them via the Submodule. Short-circuit fakes that bypass boot are a §11.4 violation. A passing test **MUST** imply the infra was up.

Tracker rows touching containerization **MUST** record

Catalogue-Check: extend `vasic-digital/containers@<sha>` (or reuse); no-match requires demonstrating the Submodule cannot model the workload.

Planned anti-bluff gate `CM-CONTAINERS-USED` scans container-touching PRs for `digital.vasic.containers/... imports`. Paired mutation strips the `import` + asserts FAIL.

Composes with §11.4.74 (catalogue-first), §11.4.75 (mechanical enforcement), §3 (propagation), §11.4.31 (Submodule-Dependency-Manifest), §11.4.36 (`install_upstreams` on clone), §11.4.28 (Submodules-As-Equal-Codebase), §1.1 (paired-mutation gates).

Canonical authority: constitution submodule Constitution.md §11.4.76.

Non-compliance: reinventing `compose` orchestration in-project is a release blocker.

§11.4.77 — Regeneration-mechanism-required mandate (User mandate, 2026-05-20)

Direct user mandate (verbatim): “We must be sure that after excluding anything from Git versioning we still have the mechanism which will out of the box obtain or re-generate missing content!”

Forensic anchor. 2026-05-20T15:00Z: ATMOSphere parent’s audio Tier 1 `commit_all.sh` stalled 4 h on `git add -A` scanning 274 GiB `.git-backup-*` + 159 GiB `RKTools/linux/` + 167 GiB `qa-results/` — all untracked but un-gitignored. Bare `.gitignore` fix would orphan every fresh clone (missing `RKTools`, missing test infra, missing build outputs).

Every `.gitignore` entry excluding (a) >~100 MiB OR (b) any artefact essential to building / running / testing the project **MUST** carry a documented + automated mechanism to either **re-obtain** (vendor tarball, SDK installer, package registry, git submodule, object store) OR **re-generate** (build pipeline, code-gen, asset render, captured-evidence replay, container build).

Required artefacts per qualifying `.gitignore` entry: (1) `.gitignore-meta/<entry-slug>.yaml` declaring pattern + mechanism-type + script-path + expected-disk-usage + vendor-url-or-source + integrity hash + requires-network + requires-credentials; (2) post-clone bootstrap entry (`scripts/setup.sh` or canonical equivalent) running the mechanism non-interactively; (3) pre-build gate verifying regenerated content present OR stamp `.gitignore-meta/.regenerated/<slug>.ok` recent; (4) `README` + `docs/guides/*.md` describing the mechanism + manual fallback + time/disk budget + per-§11.4.10 credentials.

No escape hatch: bare `.gitignore` additions without the mechanism are §11.4 PASS-bluff variants — codebase appears complete but fresh clone cannot build / run. No `--skip-regen-mechanism`, `--gitignore-is-enough`, `--operator-already-has-content` flag.

Planned anti-bluff gate `CM-GITIGNORE-REGEN-MECHANISM` scans every `.gitignore` addition for matching `.gitignore-meta/` sibling. Paired §1.1 mutation strips a required YAML key → gate FAILs.

Composes with §11.4.6 (no-guessing — verify mechanism on sandbox), §11.4.65 (HTML/PDF siblings are a re-generate instance), §11.4.66 (interactive clarification if mechanism unknown), §11.4.71 (re-validate vendor integrity before push), §11.4.74 (extend a reusable downloader Submodule), §11.4.75 (pre-commit + CI replay refuse bare additions), §11.4.76 (container images regenerated via `vasic-digital/containers`), §9 / §9.2 (pre-test mechanism in sandbox), §3 (consuming submodules inherit).

Canonical authority: constitution submodule [`Constitution.md`](#) §11.4.77.

Non-compliance is a release blocker regardless of context.

§11.4.78 — CodeGraph code-intelligence mandate (User mandate, 2026-05-20)

Direct user mandate (verbatim): “Make codegraph **MANDATORY CHOICE** for this purpose for all of our project ... All project which do not have configured and installed codegraph yet **MUST DO IT** and **MUST USE IT!**”

Every consuming project worked on by AI coding agents MUST install, initialize, and use **CodeGraph** (<https://github.com/colbymchenry/codegraph>, npm @colbymchenry/codegraph) — a local SQLite semantic code-knowledge-graph exposed to agents over MCP (100% local, no cloud). (1) Install globally via npm — no sudo (npm prefix MUST be user-writable). (2) `codegraph init + codegraph index: .codegraph/config.json` tracked, `.codegraph/codegraph.db` gitignored with `codegraph index` as its §11.4.77 regeneration mechanism; the `exclude` list MUST exclude other-owned submodules AND every §11.4.10 credential/secret path. (3) Wire the `codegraph serve --mcp` server into every CLI agent the developers use — project-scoped + committed where supported (Claude Code `.mcp.json`, OpenCode `opencode.json`, Qwen Code `.qwen/settings.json`, Crush `.crush.json`), host-local otherwise (Kimi CLI `~/.kimi/mcp.json`); configs reference the bare `codegraph` command on PATH. (4) Cover the integration with an anti-bluff verification suite whose per-agent end-to-end layer uses an unforgeable challenge (a fact obtainable only by calling a CodeGraph MCP tool); un-runnable agents are documented SKIP gaps per §11.4.3, never faked PASSES. (5) Document everything in `docs/CODEGRAPH.md`. CodeGraph is consumed as the npm package (§11.4.74) — not a git submodule, adds no Git remote.

Composes with §11.4.3, §11.4.10, §11.4.12, §11.4.65, §11.4.30, §11.4.74, §11.4.77, §11.4, §1.1. Planned gate CM-CODEGRAPH-WIRED + paired mutation.

Canonical authority: constitution submodule [Constitution.md](#) §11.4.78.

Non-compliance is a process violation; a project worked on by AI agents without CodeGraph installed, wired, and anti-bluff-verified is in breach of this mandate.

§11.4.79 — Own-org submodules MUST be included in the CodeGraph index (User mandate, 2026-05-21)

Direct user mandate (verbatim, 2026-05-21): “All Submodules we use in the project and that are part of organizations to which we have the full access via GitHub, GitLab and other CLIs MUST BE included into the codegraph database and initialized / scanned / synced!”

Refines §11.4.78 step 2’s `exclude`-list with a per-submodule-ownership split: **own-org** (full write access via project’s CLIs — canonical orgs `vasic-digital` + `HelixDevelopment`) MUST be INCLUDED; **third-party** (§11.4.74 `no-match` → `vendor`; e.g. `gopkg.in/telebot.v3`) MUST be EXCLUDED. Operational steps: (1) `git submodule update --remote --merge` before re-indexing — respect load-bearing third-party pins (roll back if `--remote` breaks a load-bearing import-path pin). (2) Adjust `.codegraph/config.json` `exclude` list. (3) Re-index. (4) Validate via probe resolving a symbol living ONLY inside an own-org submodule. (5) Paired §1.1 mutation: temporarily add own-org path to `exclude` → validate FAILs → restore.

Composes with §11.4.74 (ownership classifier), §11.4.78 (CodeGraph parent), §11.4.10 (credentials still excluded), §1.1 (paired mutation), §107 (a lying index is a PASS-bluff).

Canonical authority: constitution submodule [Constitution.md](#) §11.4.79.

Non-compliance is a process violation; severe cases (own-org submodules silently excluded without an audit trail in `.codegraph/config.json`) are release blockers.

§11.4.80 — CodeGraph regular-update + sync automation mandate (User mandate, 2026-05-21)

Direct user mandate (verbatim, 2026-05-21): “We MUST regularly check for the updates and execute codegraph npm updates so the latest version of it is always installed on the host machine! [...] Make sure we have proper full automation bash scripts which will run regularly and that these are part of the constitution Submodule since they MUST BE available to all projects [...]. Make sure all updates, sync processes we do and important codegraph related events are all documented under docs/codegraph in Status and Status_Summary documents [...]”

Three deliverables in this constitution submodule: (1) scripts/codegraph_update.sh — npm-installs latest; §107 anti-bluff verifies codegraph --version reflects new version (npm exit 0 ≠ working binary). (2) scripts/codegraph_sync.sh — runs codegraph status → sync . → status → project’s validate; appends to both project’s + constitution’s docs/codegraph/Status.md. (3) docs/codegraph/Status.md + Status_Summary.md ledgers, exported to .html + .pdf per §11.4.65.

Cadence: weekly floor. Scripts inherited by reference (§3 submodule inheritance) — invoke at \${CONST_DIR}/scripts/codegraph_*.sh, never copy.

Composes with §11.4.78, §11.4.79, §11.4.10, §11.4.45, §11.4.65, §107, §1.1.

Canonical authority: constitution submodule Constitution.md §11.4.80.

Non-compliance is a process violation; severe cases (>2 weeks since last update + open AI-agent work) are release blockers.

§11.4.81 — Cross-platform-parity mandate (User mandate, 2026-05-21)

Direct user mandate (verbatim, 2026-05-21): “Any Linux-only blocker / issue we have MUST BE created macOS and other supported platforms equivalent! So, depending on platform proper implementation will be used for particular OS! EVERYTHING MUST BE PROPERLY EXTENDED AND UPDATED!”

Every consuming project whose supported-platforms manifest lists more than one OS MUST ship per-OS-equivalent implementations + tests for every feature/gate/challenge/mutation that depends on platform-specific primitives. Runtime dispatch via uname -s (or equivalent). Three sub-mandates: **(A)** Per-OS implementation REQUIRED — cgroup/systemd/proc primitives have documented equivalents (POSIX setrlimit/ulimit, launchd, rctl, Job Object). **(B)** Per-OS tests REQUIRED — every gate test has case "\$(uname -s)" in branches with positive captured evidence per §11.4.2 + §11.4.5 in each branch; SKIP-with-reason only when the platform genuinely cannot enforce. **(C)** Honest kernel-gap citation + adjacent equivalent test REQUIRED where no equivalent exists (canonical: XNU does NOT enforce RLIMIT_AS for unprivileged processes → SKIP with exact reproducer + adjacent test of what IS enforced, e.g. RLIMIT_CPU+SIGXCPU). The adjacent test is itself anti-bluff per §11.4 with a paired §1.1 mutation.

Per-OS equivalence catalogue (canonical): systemd-run --user --scope [?] POSIX ulimit -t -u / launchd; cgroup MemoryMax [?] XNU gap (use RLIMIT_CPU adjacent) / rctl / Job Object; cgroup TasksMax [?] RLIMIT_NPROC; / proc/<pid>/oom_score_adj [?] no Darwin/BSD equivalent.

Composes with §11.4.1 / §11.4.2 / §11.4.3 (strictened — SKIP only when kernel cannot) / §11.4.4 / §11.4.5 / §11.4.6 / §11.4.20 / §11.4.27 / §11.4.69 / §11.4.70 / §107. Pre-build gate CM-CROSS-PLATFORM-PARITY + paired §1.1 mutation. No escape hatch.

Canonical authority: constitution submodule Constitution.md §11.4.81.

Non-compliance is a release blocker on multi-platform projects.

§11.4.82 — Iteration-speedup discipline mandate (User mandate, 2026-05-22)

Direct user mandate (verbatim, 2026-05-22): “How can we speed-up this whole development and fixing process? ... all speed optimizations critical rules and mandatory constraints MUST BE all added into our root (constitution Submodule) Constitution.md, CLAUDE.md, AGENTS.md and QWEN.md!”

Iteration cycle time bounds the project’s defect-discovery rate. Slow cycles ship fewer validated features per unit of operator time. Every consuming project’s build / test / commit / debug pipeline MUST adopt 9 speedup disciplines AS MANDATORY:

- **(A) Phase 1 forensic before any speculative source patch** — Phase 1 of superpowers:systematic-debugging is mandatory before any non-trivial source patch. Speculative patches violate §11.4.6 + §11.4.82.
- **(B) Live-ADB-First before rebuild** — §11.4.51 strengthened to release-blocker. Skipping live-probe for LIVE_ADB_TESTABLE changes = §11.4.82 violation.
- **(C) Pre-flight before rebuild orchestrators** — 30-sec pre-flight: device reachable, sinks reachable, memory budget, no stale locks, no orphan processes.
- **(D) Persistent build caches outside containers** — ccache + Soong / sccache / Gradle daemon state bind-mounted to host. Saves 25-32 min/cycle from cycle #2.
- **(E) Module-only rebuild for CONFIG_*=m driver patches** — make modules over single driver dir (~2-5 min vs ~25 min full kernel).
- **(F) Parallel multi-device testing** — every owned validation device runs autonomous cycle concurrently with separate output dirs.
- **(G) Subagent scope ≤30 min + worktree isolation + single-responsibility** — empirical pattern: 8-task subagents stall, 2-3-task subagents complete.
- **(H) Lock-file + stale-process hygiene** — clean .git/index.lock / .lock / orphan git/build processes on session start. Disable auto git-gc in concurrent multi-agent repos.
- **(I) Cycle telemetry per §11.4.24** — log commit hash, per-phase wall-clock, speedup flag set, outcome. Weekly aggregation surfaces empirical ROI per discipline.

Composes with §11.4.4 / §11.4.6 / §11.4.9 / §11.4.20 / §11.4.24 / §11.4.42 / §11.4.43 / §11.4.50 / §11.4.51 / §11.4.52 / §11.4.58 / §11.4.70 / §12.7 / §107. Pre-build gate CM-ITERATION-SPEEDUP-DISCIPLINE (when implemented) + paired §1.1 mutation. No escape hatch — no --skip-phase1, --no-pre-flight, --rebuild-everything, --unlimited-subagent-scope, --ignore-locks, --no-telemetry flag exists.

Canonical authority: constitution submodule Constitution.md §11.4.82.

Non-compliance is a release blocker.

§11.4.83 — docs/qa/ end-user evidence mandate (User mandate, 2026-05-22)

Direct user mandate (verbatim, 2026-05-22): “every feature that ships MUST carry a recorded e2e communication transcript + any attached materials under docs/qa/<run-id>/ (per-feature subdirectories). A feature with no QA transcript is itself a §107 PASS-bluff — it claims to work but has no auditable runtime evidence. Bot-driven automation MUST preserve full bidirectional communication threads as proof.”

Every feature that ships MUST carry a recorded end-to-end communication transcript plus any attached materials committed under docs/qa/<run-id>/. A feature with no QA transcript is itself a §11.4 / §107 PASS-bluff. (1) Transcripts MUST be full bidirectional. (2) Attached materials MUST live in-repo (no external-only links — §11.4.13 sink-side violation). (3) Bot-driven / agent-driven QA MUST preserve the full conversation thread as the proof artefact. (4) CI release gates MUST refuse to tag a version whose feature-shipping commit lacks its matching docs/qa/<run-id>/.

Composes with §11.4.2, §11.4.5, §11.4.13, §11.4.65, §11.4.69, §107, §1.1.

Canonical authority: constitution submodule Constitution.md §11.4.83.

Non-compliance is a release blocker.

§11.4.84 — Working-tree quiescence rule for subagent commits (User mandate, 2026-05-22)

Short tag: working-tree quiescence.

Direct user mandate (verbatim, 2026-05-22): “no subagent commit may proceed while any concurrent mutation gate is in flight in the same checkout. Before `git add`, the committing agent MUST `grep` its own working tree for mutation markers (MUTATED for paired, // always pass, return `json.Marshal` shortcut paths, etc.). Any unexplained file in the staging area triggers ABORT.”

No subagent or main-thread commit may proceed while any concurrent mutation gate is in flight in the same checkout. Pre-`git add`: `grep` for mutation markers (MUTATED for paired, // always pass, return `json.Marshal` shortcuts, // MUTATION / # MUTATION annotations, `_mutated_*` filenames). Cross-check `git status --porcelain` against declared scope — unaccounted entries → ABORT.

Lesson (forensic). A consuming project’s logo-fix subagent (Herald commit 72e81ab, 2026-05-21) ran in a checkout where a paired §1.1 mutation had injected an // always pass shortcut into a JWT verify path. The subagent’s `git add` swept the mutation residue into the commit; the resulting commit was pushed to all four mirrors before any other agent caught it. Fix (Herald d5bd360) landed within the hour, but the production-equivalent-binary-with-bypassed-JWT window is a real security-defect window. The lesson is now constitutional.

Operative rule. (1) Pre-`git add` `grep` + scope-match → unaccounted file ABORT. (2) Active mutation gates MUST be serialised — mutate → assert FAIL → restore → assert PASS — and tree clean BEFORE unrelated commits. (3) Concurrent subagents same-checkout MUST use a lockfile (`.git/MUTATION_IN_PROGRESS`) or `git worktree add` per subagent. (4) Pre-push mutation-residue-scanner MUST run; commits containing mutation markers → push BLOCKED.

Composes with §1.1, §11.4.20, §11.4.70, §11.4.27, §11.4.10, §11.4.71, §107.

Canonical authority: constitution submodule Constitution.md §11.4.84.

Non-compliance is a release blocker. A mutation marker in a tagged commit is a critical defect regardless of how briefly it persisted.

§11.4.85 — Stress + Chaos Test Mandate (User mandate, 2026-05-24)

Short tag: stress-chaos-mandate.

Direct user mandate (verbatim, 2026-05-24): “Every fix or improvement you do MUST BE covered with full automation stress and chaos tests so we are sure nothing can break the functionality and all edge cases are monitored and polished and additionally fixed if that is needed! Everything must produce rock solid proofs and follow fully no-bluff policy!”

Every fix MUST ship with full-automation stress + chaos test suites. Happy-path-only coverage is a §11.4 / §107 PASS-bluff at the resilience layer.

Stress (closed-set): sustained load ($N \geq 100$ iters OR ≥ 30 s, p50/p95/p99 recorded) + concurrent contention ($N \geq 10$ parallel, no deadlock, no leak) + boundary conditions (empty/max/off-by-one, every boundary produces categorised result).

Chaos (closed-set, per fix-class): process-death injection + network-fault injection + input-corruption injection + resource-exhaustion injection (disk/OOM/FD) + state-corruption injection. Recovery deterministic + categorised per §11.4.69.

Anti-bluff. Every stress + chaos PASS cites a captured-evidence artefact (latency.json / categorised_errors.txt / state_delta_snapshot.json / recovery_trace.log) per §11.4.5 + §11.4.69. Helper library stress_chaos.sh provides ab_stress_run / ab_stress_concurrent / ab_chaos_* primitives composing with ab_pass_with_evidence / ab_skip_with_reason. Chaos cleanup MUST run in trap EXIT — failure = §11.4.14 violation.

4-layer per §11.4.4(b): pre-build gate + paired meta-test mutation + on-device test (if LIVE_ADB_TESTABLE) + HelixQA Challenge entry. Composes with §11.4 / §11.4.1 / §11.4.5 / §11.4.6 / §11.4.43 / §11.4.50 / §11.4.52 / §11.4.69 / §11.4.83 / §107.

Canonical authority: constitution submodule Constitution.md §11.4.85.

Non-compliance is a release blocker regardless of context. No --skip-stress, --no-chaos, --happy-path-suffices, --stress-test-later flag exists.

§11.4.86 — Roster/corpus-backed Status-doc auto-sync mandate (User mandate, 2026-05-25)

Forensic anchor (verbatim): “Make sure that assets and players Status docs are ALWAYS regularly updated and in sync like all others Status docs — any time we add or modify the assets content(s) or we change or add new / remove existing pre-installed video and audio player apps! This MUST WORK OUT OF THE BOX!”

Status docs (§11.4.45) backed by a **tracked roster** (installed apps/components) or **asset corpus** (test/media asset directory) MUST resync **out of the box** the moment a member changes (player added/removed/renamed; asset added/modified/removed) — Status doc + Status_Summary + HTML + PDF, mechanically, not by operator vigilance.

Mechanism (all must hold): (1) drift-proof **fingerprint** = sha256 of sorted member list (NOT mtime), persisted in a sidecar beside the Status doc; (2) **sync helper** regenerates fingerprint + re-exports HTML+PDF (via §11.4.65 exporter), wired so sync is automatic; (3) **pre-build gate** FAILs when live fingerprint ≠ persisted (mirrors §11.4.12 CM-ISSUES-SUMMARY-SYNC + §11.4.45 sync_integration_status); (4) paired §1.1 mutation corrupts fingerprint → gate FAILs. Composes with §11.4.12 / §11.4.45 / §11.4.53 / §11.4.56 / §11.4.57 / §11.4.59 / §11.4.60 / §11.4.65 / §11.4.6. Classification universal (§11.4.17) — consuming project supplies specific docs/roster/helper/gate per §11.4.35.

Canonical authority: constitution submodule Constitution.md §11.4.86.

Non-compliance is a release blocker regardless of context. No --skip-roster-sync, --allow-status-drift, --roster-sync-not-applicable flag exists.

§11.4.87 — Endless-loop autonomous work + zero-idle agent dispatch + anti-bluff testing mandate (User mandate, 2026-05-26)

When operator instructs an AI agent to “continue in endless loop fully autonomously” (or semantically-equivalent), the agent MUST treat as HARD-CONTRACT covenant covering five obligations: (A) continue until docs/Issues.md non-terminal Status entries = 0 AND docs/CONTINUATION.md §3 Active work empty AND no subagent in-flight AND no external dep in-flight; (B) dispatch background subagents for parallelisable work — main + subagents concurrent, “waiting for results” is the ONLY idle reason; (C) every closure lands four-layer test coverage per §11.4.4(b) with captured-evidence “physical proofs” (tinycap WAV + RMS / screen recording + fprobe / dumphsys + sink-probe / uiautomator dump / sysfs snapshots) — metadata-only / config-only / absence-of-error / grep-without-runtime PASS are critical defects; (D) §11.4 anti-bluff covenant family operative end-to-end (tests AND HelixQA Challenges bound equally per forensic anchor “tests pass but features don’t work”); (E) loop terminates ONLY on all-conditions-met, explicit operator STOP, host-safety demand (§12 family), or scheduled wake on known-future-actionable signal.

Composes with §11.4 / §11.4.1 / §11.4.2 / §11.4.4 / §11.4.5 / §11.4.6 / §11.4.7 / §11.4.20 / §11.4.27 / §11.4.42 / §11.4.43 / §11.4.50 / §11.4.52 / §11.4.58 / §11.4.68 / §11.4.69 / §11.4.70 / §11.4.83 / §11.4.85 / §11.4.86 / §12.10. Pre-build gate CM-COVENANT-114-87-PROPAGATION + paired §1.1 mutation.

Canonical authority: constitution submodule Constitution.md §11.4.87.

Non-compliance is a release blocker regardless of context. No escape hatch — --idle-OK, --skip-endless-loop, --bluff-permitted-for-this-task, --metadata-only-test-suffices, --no-physical-proof-required are FORBIDDEN flags.

§11.4.88 — Background-push mandate: commit-lock release immediately after commit, push runs detached (User mandate, 2026-05-26)

Forensic anchor: 2026-05-26 a single `commit_all.sh` held its flock ~5 hours because `do_push` ran synchronously after the commit landed — every subsequent commit was blocked on a slow mirror push that had nothing to do with the local commit's durability. Implementation seam for §11.4.87(B) zero-idle.

Mandate (5 obligations): (A) `.git/.commit_all.lock` released IMMEDIATELY after `git commit` returns 0; (B) push runs detached via `nohub ./push_all.sh ... > <log> 2>&1 & + disown`; (C) `push_all.sh` acquires per-remote flock `.git/.push.<remote>.lock` — concurrent same-remote serializes, different-remote parallelises; (D) failures → `qa-results/push_failures/<ts>_<remote>.log`, autonomous loop checks per §11.4.87(A); (E) `--sync-push` flag escape for §11.4.41 force-push merge-first paths only. Composes with §2.1 / §9.2 / §11.4.41 / §11.4.42 / §11.4.71 / §11.4.87. Pre-build gates CM-COVENANT-114-88-PROPAGATION + CM-BACKGROUND-PUSH-WIRED + paired §1.1 mutations.

Canonical authority: constitution submodule Constitution.md §11.4.88.

Non-compliance is a release blocker. No escape hatch beyond `--sync-push`.

§11.4.89 — Background test execution mandate (User mandate, 2026-05-27)

Forensic anchor — verbatim user mandate (2026-05-27): “Any tests we are executing, especially long test cycles, MUST BE performed in background in parallel with main work stream! This MUST NOT block our capabilities to work on queued workable items.”

Symmetric anchor to §11.4.88 at the test-execution layer. Mandate (6 obligations): (A) Long-running tests (>30s) MUST run via `nohub ... > <log> 2>&1 & + disown` with log under known dir; (B) Main stream proceeds to §11.4.42 priority queue immediately — “waiting for results” is the only acceptable idle reason per §11.4.87(B); (C) Hard-dependency gating: poll exit-status file or `pgrep` before dependent steps — surface as §11.4.66 options if test still running; (D) Failures land in log files; (E) Foreground only for <30s OR operator authorisation; (F) Per-script flock for same-script serialisation. Composes with §11.4.42 / §11.4.66 / §11.4.82 / §11.4.84 / §11.4.85 / §11.4.87 / §11.4.88. Pre-build gates CM-COVENANT-114-89-PROPAGATION + CM-BACKGROUND-TEST-EXECUTION-WIRED + paired §1.1 mutations.

Canonical authority: constitution submodule Constitution.md §11.4.89.

Non-compliance is a release blocker. No escape hatch beyond explicit per-invocation operator authorisation.

§11.4.90 — Obsolete status + per-item obsolescence audit (User mandate, 2026-05-27)

§11.4.15 Status closed-set extended with terminal `Obsolete` (→ `Fixed.md`) (orthogonal to `Type`). Reasons (closed): `superseded-by-design-change` | `superseded-by-later-mandate` | `feature-removed` | `duplicate-of` | `unsupported-topology` | `not-reproducible` (`not-reproducible` = a reported defect that does NOT reproduce on the canonical tree/baseline — an environment/

isolated-worktree artifact, not a real product defect). Every Obsolete heading carries ****Obsolete-Details:**** (Since + Reason + Superseding-item + Triple-check evidence). Colorizer cell-status-obsolete = #E0E0E0 background + strikethrough. Audit cadence per §11.4.40 + §11.4.42. Triple-check non-negotiable. Pre-build gates CM-COVENANT-114-90-PROPAGATION + CM-ITEM-OBSOLETE-DETAILS + CM-OBSOLETE-COLORIZER-WIRED + paired §1.1 mutations.

Canonical authority: constitution submodule Constitution.md §11.4.90. Non-compliance is a release blocker.

§11.4.91 — Summary-doc clarity mandate (User mandate, 2026-05-27)

Every summary entry MUST contain self-contained meaningful description ≥ 6 words OR ≥ 40 chars naming SUBJECT + PROBLEM/GOAL. Forbidden anti-patterns: section labels (Composes with, etc.); bare metadata fragments; §-letter alone. Generators MUST extract from H1/H2 heading line per §11.4.54 ATM-NNN convention. Pre-build gate CM-SUMMARY-CLARITY-DESCRIPTIONS + paired §1.1 mutation.

Canonical authority: constitution submodule Constitution.md §11.4.91. Non-compliance is a release blocker.

§11.4.92 — Multi-pass change-evaluation discipline (User mandate, 2026-05-27)

Every non-trivial change MUST pass 5-pass evaluation BEFORE commit-ready: (1) Main-task verification, (2) Regression-blast-radius, (3) Cross-feature interaction, (4) Deep-research validation per §11.4.8, (5) Anti-bluff confirmation. Documentation per pass. Only after all 5 passes complete may commit/push/test/release proceed. Composes §11.4.4 / §11.4.5 / §11.4.6 / §11.4.8 / §11.4.20 / §11.4.27 / §11.4.42 / §11.4.43 / §11.4.50 / §11.4.52 / §11.4.69 / §11.4.78 / §11.4.82 / §11.4.83 / §11.4.85 / §11.4.87 / §11.4.89. Pre-build gates CM-COVENANT-114-92-PROPAGATION + CM-MULTI-PASS-EVALUATION-EVIDENCE.

Canonical authority: constitution submodule Constitution.md §11.4.92. Non-compliance is a release blocker.

§11.4.93 — SQLite-backed single-source-of-truth for workable items (User mandate, 2026-05-27)

Text-based Issues/Fixed/Summary trackers converted to SQLite DB at docs/.workable_items.db (gitignored + §11.4.77 regen mechanism). Schema: items (atm_id PK + Type + Status incl. Obsolete + Severity + title + description ≥ 40 chars + composes_with JSON + current_location); item_history (audit log §11.4.34); obsolete_details (§11.4.90); operator_block_details (§11.4.21); firebase_metadata (§11.4.47); meta. Go binary cmd/workable-items/: sync md-to-db / db-to-md / diff / validate / add / close. Bidirectional byte-identical round-trip. commit_all.sh refuses on diff. Pre-build gate runs validate. Anti-bluff: unit + integration + stress + chaos per §11.4.85 + paired §1.1 mutation + HelixQA

Challenge CME-WORKABLE-ITEMS-001. 6-phase migration §LA. Go binary in constitution submodule per §11.4.74. Composes §11.4 / §11.4.12-93 covenant family.

Canonical authority: constitution submodule Constitution.md §11.4.93. Non-compliance is a release blocker.

§11.4.94 — Zero-idle priority-first parallel-by-default operating mode (User mandate, 2026-05-27)

Binds

§11.4.20+§11.4.42+§11.4.58+§11.4.70+§11.4.72+§11.4.82+§11.4.87+§11.4.88+§11.4.89 into single always-on enforcement: idle only when externally blocked / operator STOP / §12 host-safety; before any wake/sleep survey parallel-work feasibility; priority order MANDATORY (audio-first per §11.4.72); subagent-driven default non-trivial; background default >30s; stability-preserving per §11.4.92 + §11.4.84 + §12.6-§12.9; progress updates surfaced at milestones. Anti-pattern: wake without queue survey, “nothing to do” with non-trivial progressable, serial parallelisable work, lower priority over higher. Pre-build gates CM-COVENANT-114-94-PROPAGATION + CM-PARALLEL-WORK-AUDIT.

Canonical authority: constitution submodule Constitution.md §11.4.94. Non-compliance is a release blocker.

§11.4.95 — Workable-items SQLite DB TRACKED in git, NEVER gitignored (User mandate, 2026-05-27)

§11.4.93 amendment: DB at docs/workable_items.db is TRACKED, NEVER gitignored. Authoritative source data, NOT build artefact. Every state change → stage+commit+push DB alongside MD regen per §11.4.19+§2.1. WAL-checkpoint before commit. Destructive DB ops require §9.2+authorization. Pre-build gates CM-COVENANT-114-95-PROPAGATION + CM-WORKABLE-ITEMS-DB-TRACKED.

Canonical authority: constitution submodule Constitution.md §11.4.95. Non-compliance is a release blocker.

§11.4.96 — Safe-parallel-work-with-long-build catalogue + mandate (User mandate, 2026-05-27)

Operational catalogue for 5-7h AOSP containerised build:

SAFE: (A) docs/MD work; (B) host-side scripts/; (C) pre-build + meta-test gate authoring; (D) on-device test scripts authoring; (E) constitution edits + push; (F) any submodule commit + push per §11.4.88; (G) read-only live-ADB probes; (H) subagent dispatch per §11.4.20/§11.4.70 + §11.4.84; (I) web research + APIs; (J) workable-items DB ops; (K) pre-build + meta-test execution backgrounded per §11.4.89.

UNSAFE: (α) git checkout/reset on source tree; (β) mass deletes/renames under built source; (γ) submodule pointer updates affecting built APKs; (δ) out/ mutations; (ε) make clean / m clobber; (ζ) container destruction; (η) disk-filling; (θ) §12 host-safety breaches.

Conductor MUST dispatch ≥ 1 (A)-(K) item per pause point. “Build running, nothing else to do” NEVER true. Pre-build gates CM-COVENANT-114-96-PROPAGATION + CM-PARALLEL-WORK-DURING-BUILD-AUDIT.

Canonical authority: constitution submodule Constitution.md §11.4.96. Non-compliance is a release blocker.

§11.4.97 — Maximum-use-of-idle-time + progress-update cadence (User mandate, 2026-05-27)

Operating-mode capstone: (A) every minute of progressable idle time = §11.4.97 violation; dispatch CONTINUOUSLY; (B) emit 1-line operator update at every commit/subagent-return/anchor-landed/evidence-captured/migration-closure — no prompt required; (C) continuous physical-proof gathering per §11.4.5+§11.4.6+§11.4.69; (D) composes with §11.4.5/6/13/20/27/42/50/52/69/70/72/83/85/87/88/89/94/96; (E) idle-only-when-blocked closed-set unchanged. Pre-build gates CM-COVENANT-114-97-PROPAGATION + CM-IDLE-TIME-AUDIT.

Canonical authority: constitution submodule Constitution.md §11.4.97. Non-compliance is a release blocker.

§11.4.98 — Full-Automation Anti-Bluff Mandate — Live tests MUST be re-runnable end-to-end without manual intervention (User mandate, 2026-05-28)

Forensic anchor — verbatim user mandate (2026-05-28): “Make sure we have full automation testing of all scenarios with real bot, main group and users without any manual intervention or contribution of real user! ... No bluff is allowed!”

Closes the manual-intervention gap §11.4 + §11.4.85 + §11.4.87 + §11.4.89 + §11.4.94 did not explicitly forbid: a test requiring human action during execution is a §11.4 PASS-bluff at the automation layer. (A) Binding rule: every test MUST be self-driving end-to-end; PASS/FAIL/SKIP-with-reason without further human action after startup. (B) Single exception: one-time credential bootstrap OUTSIDE test execution (.env, ~/.bashrc, OAuth, MTPROTO-session-activation) — configuration, not test driving. (C) Concrete live-messenger requirements: (1) no “operator MUST type” prompts — drive programmatically (MTPROTO/real-user-API/webhook-fixture/in-process loopback); (2) no UUIDs colliding with active dev session (Herald 2026-05-28 lesson: same-UUID claude --resume returns silent exit -1); (3) no 60s human-response windows (§11.4.50 determinism violation); (4) re-runnability proof at -count=3 consecutive automated runs with self-cleaning state; (5) §11.4.98 audit — every existing test classified COMPLIANT vs NON-COMPLIANT; (6) no false-positive PASS — silent-skip-as-PASS forbidden, stale-evidence forbidden, SKIP-with-reason per §11.4.3. (D) Composes with §11.4.85+89+87+94 = continuously-validated fully-automated non-flake anti-bluff regime. (E) Inheritance per §11.4.35 — restate literal anchor 11.4.98 in every consumer’s CLAUDE/AGENTS/QWEN; pre-build gate CM-COVENANT-114-98-PROPAGATION; paired §1.1 mutations strip → gates FAIL. (F) Enforcement: manual-action commit BLOCKED at release-gate; NON-COMPLIANT test not rewritten in 30 days → §11.4.90 Obsolete citing §11.4.98.

Canonical authority: constitution submodule Constitution.md §11.4.98. Non-compliance is a release blocker.

§11.4.99 — Latest-Source Documentation Cross-Reference Mandate — instructions/guides/manuals MUST be verified against latest official online sources BEFORE publication (User mandate, 2026-05-28)

Forensic anchor: “Make sure we ALWAYS check against latest versions of services we use web/online docs before creating instructions! This situation is illustration of how we can misguide ourselves or get banned!”

Case study (Herald 2026-05-28): first-draft MTPProto guide recommended VoIP + omitted `recover@telegram.org` pre-login email; both contradicted official Telegram + gotd/td docs and could have caused permanent account ban. Misguidance-by-stale-docs = §11.4 PASS-bluff at documentation layer.

1. Pre-commit cross-reference: agent MUST fetch LATEST official online docs (WebFetch/MCP/direct browsing — NEVER training data) for every operator-facing instruction doc; cross-reference each instruction; seek secondary authoritative sources (maintainer SUPPORT.md, changelogs, vetted FAQs) when official is sparse; cite source URL + date in “## Sources verified” footer; cite in commit-message footer (Sources verified <date>: <urls>). (B) Negative findings documented (gaps/silences = no authoritative agreement). (C) Re-verification cadence — STALE after 6 months (90 days for risk-classified services); re-verify before operator-authority citation, at vN.0.0 release boundary, on breaking-change announcements, on operator-error report. (D) Risk-classified services (messaging/cloud/payment/AI/code-hosting/package-managers) — 90-day staleness limit + explicit safety warnings vs latest policies. (E) Composes with §11.4.92 Pass 4 but INDEPENDENT — cannot substitute. (F) Inheritance per §11.4.35 — restate literal 11.4.99 in every consumer; pre-build gate CM-COVENANT-114-99-PROPAGATION. (G) Enforcement: missing footer → BLOCKED at release-gate; stale-beyond-grace → §11.4.90 Obsolete Reason=stale-documentation.

Canonical authority: constitution submodule Constitution.md §11.4.99. Non-compliance is a release blocker.

§11.4.100 — RETIRED. Demoted to consumer project (ATMOSphere video-color/visual-quality fidelity) per §11.4.17/§11.4.35 — project-specific (RK3588/MPV/Arvus), not universal. See the consuming project’s Constitution/CLAUDE/AGENTS/QWEN.

§11.4.101 — Autonomous-decision-over-blocking mandate (User mandate, 2026-05-28)

Forensic anchor — verbatim user mandate (2026-05-28): “when working in endless working loop fully autonomously try to decide most properly about points which would block execution and wait for us. If we haven’t answered now work would be blocked whole night! If possible and if that will not cause any issues make proper and most reliable and safe decision so we achieve maximal efficiency and work gets fully done!”

When operating in autonomous / endless-loop mode (per §11.4.87), the agent **MUST** minimize operator-blocking and instead make the safe, reliable, reversible decision itself — so work is NOT stalled (e.g. overnight) waiting for input. §11.4.87 says keep working; §11.4.101 says HOW to clear the decision points.

Decision rule (closed-set — proceed autonomously when ALL hold): (a) action reversible OR captured pre-op backup per §9.2; (b) safe choice determinable from captured evidence per §11.4.6 (no guessing — LIKELY is not a determination); (c) wrong-choice blast radius bounded AND recoverable; (d) composes with anti-bluff §11.4 + host-safety §12 + data-safety §9.

Block-only-when rule (BLOCK via §11.4.66 ONLY when ALL hold): action irreversible AND high-blast-radius AND safe choice undeterminable from evidence — e.g. external-account state the agent cannot inspect, hardware it cannot access, destructive ops without backup, force-push (also §9.2 + §11.4.41), spending / sending to third parties. Operator-blocked per §11.4.21 reached only after this rule fires AND self-resolution-exhaustion audit completes.

Maximize-progress-while-blocked: an unavoidable block parks one work unit, not the loop — keep progressing every NON-blocked item in parallel per §11.4.87 + §11.4.94. Posing the question and going idle is a §11.4.94 + §11.4.97 violation.

Composes with §11.4.6 / §11.4.21 / §11.4.40 / §11.4.41 / §11.4.66 / §11.4.87 / §11.4.94 / §9.2 / §12. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-101-PROPAGATION (literal 11.4.101 across consumer fleet) + paired §1.1 meta-test mutation (strip literal → gate FAILs; gate-code = separate work item). No escape hatch — no --always-block-on-decision, --never-decide-autonomously, --skip-decision-rule, --block-without-self-resolution flag.

Canonical authority: constitution submodule Constitution.md §11.4.101. Non-compliance is a release blocker.

§11.4.102 — Mandatory systematic-debugging activation + always-loaded skill-discovery + plugin-dependency availability (User mandate, 2026-05-29)

Forensic anchor — verbatim user mandate (2026-05-29): “Make sure that we ALWAYS trigger / start the”/superpowers:systematic-debugging” skills when any issues happen! If this is possible to activate and use in this situations out of the box when we spot problems / issues / bugs / misalignments / inconsistencies we **MUST** activate the skill(s) and make strongest efforts in full in depth analysis / debugging and determine root causes of all problem or obtain relevant data and information we need! ... we **MUST** make sure that “/using-superpowers” skill is ALWAYS loaded, applied and used! All dependencies (plugins) that Claude Code or other market places are offering **MUST BE** installed if these are not already available for loading and use!”

Three cooperating invariants — the difference between guess-and-retry and investigate-to-root-cause-first.

(A) Mandatory systematic-debugging activation. On ANY spotted issue / bug / test failure / gate failure / regression / misalignment / inconsistency / unexpected behaviour, the agent **MUST** activate superpowers:systematic-debugging (or

platform-equivalent) BEFORE proposing/writing/applying any fix — Iron Law: NO FIXES WITHOUT ROOT CAUSE INVESTIGATION FIRST. Full four-phase arc: root-cause (read real failure, reproduce, gather facts) → pattern (classify defect, search same pattern elsewhere) → hypothesis (falsifiable cause, prove/disprove with captured evidence) → implementation (fix only against proven root cause). Guess-and-retry, symptom-patching, re-running a failed test hoping it passes (“probably transient/flaky”) WITHOUT a completed investigation are §11.4.102 violations. Real-world signal: calling a failure transient/flaky/intermittent/probably-timing without captured forensic evidence is simultaneously a §11.4.6 (no-guessing) + §11.4.7 (demotion-evidence) violation — §11.4.102 makes the corrective response mechanical.

(B) Mandatory always-loaded using-superpowers. superpowers:using-superpowers (or platform-equivalent skill-discovery discipline) MUST be loaded + applied at session start and consulted before any task: survey available skills before acting; if ANY skill could apply — even at 1% relevance — it MUST be invoked rather than improvised. Skipping skill-discovery, or improvising a task a loaded skill exists for, is a §11.4.102 violation.

(C) Mandatory plugin / dependency availability. Every skill plugin / marketplace package / capability dependency the project relies on (Claude Code, its marketplaces, or any other runtime’s plugin ecosystem) MUST be installed + loadable BEFORE dependent work proceeds. A missing plugin blocking a mandated skill is a release-blocker until installed + confirmed loadable. Install mechanism: the runtime’s own plugin/marketplace path — for Claude Code the in-session /plugin flow (add marketplace → install plugin → confirm skill appears in available-skills list); for other runtimes the documented package/extension installer. Anti-bluff: confirm by observing the skill in the live capability list (install exit 0 ≠ skill loadable, per §11.4.80 lesson).

Composes with §11.4.4 / §11.4.6 / §11.4.7 / §11.4.8 / §11.4.43 / §11.4.70 / §11.4.82(A) / §11.4.92. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-102-PROPAGATION (literal 11.4.102 across consumer fleet) + paired §1.1 meta-test mutation (strip literal → gate FAILs; gate-code = separate work item). No escape hatch — no --skip-systematic-debugging, --guess-and-retry-OK, --symptom-patch-permitted, --skip-skill-discovery, --plugin-optional, --missing-plugin-is-warning flag.

Canonical authority: constitution submodule Constitution.md §11.4.102. Non-compliance is a release blocker.

§11.4.103 — Continuous parallel-stream working routine (User mandate, 2026-05-29)

Forensic anchor — verbatim user mandate (2026-05-29): “Do this working approach continuously and make it part of regular working routine and add it to the Constitution Submodule documented fully.”

Promotes the proven multi-stream operating pattern into the project’s standing default working routine (not a per-request opt-in). Adds the load-bearing invariant: the main work stream MUST always stay FREE. (A) Main stream FREE — ALL commit AND push run detached (nohop ... & + disown per §11.4.88), never blocking on a push or slow mirror. (B) ≥3 parallel background streams at all times + auto-backfill (User mandate 2026-05-29; raised from ≥2) — run at least three

subagent-driven background streams (per §11.4.70/§11.4.20, isolated per §11.4.58 + §11.4.84) alongside the main stream whenever three-plus non-contending actionable items exist; the moment any one stream is FULLY done, a new stream MUST immediately start and take its place (claim next-highest-priority non-contending item), so the active-stream count NEVER drops below 3 while actionable items remain. Idle below 3 only when no remaining non-contending actionable items OR all externally blocked (§11.4.94/§11.4.97/§11.4.101). Standing band 3–6, bounded above by §12.6 60% memory + §11.4.58 6-agent cap. (C) Most-critical + most-visible first; audio always top per §11.4.72 + §11.4.42. (D) Safe-during-build scope only — while a heavy gradle/m -j/42 GB containerised AOSP build (§12.9) runs, streams restrict to the §11.4.96 SAFE catalogue; NEVER a second concurrent heavy build (§12.8). (E) Heavy anti-bluff on every closure (§11.4.5/§11.4.6/§11.4.50/§11.4.69/§11.4.102). (F) Idle ONLY when genuinely externally blocked (hardware/network/in-flight build-test-push) OR operator STOP OR §12 host-safety per §11.4.94(A)+§11.4.97; a block parks one work unit, not the loop (§11.4.101).

Composes with §11.4.58 / §11.4.70 / §11.4.72 / §11.4.87 / §11.4.88 / §11.4.89 / §11.4.94 / §11.4.96 / §11.4.97 / §11.4.101 / §11.4.102 / §11.4.42 / §11.4.84 / §12.6-§12.9 / §9.2. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-103-PROPAGATION (literal 11.4.103 across consumer fleet) + paired §1.1 meta-test mutation (strip literal → gate FAILs; gate-code = separate work item). No escape hatch — no --block-main-stream, --synchronous-commit, --synchronous-push, --single-stream-only, --skip-parallel-streams, --serialise-actionable-work, --idle-without-queue-survey flag.

Canonical authority: constitution submodule Constitution.md §11.4.103. Non-compliance is a release blocker.

§11.4.104 — Participant identity, attribution & notification-tagging (User mandate, 2026-05-31)

Forensic anchor — verbatim user mandate (2026-05-31): “Every supported messenger must relate messages to participants (Subscribers/Users); the same logical person may have a different username on every messenger. Workable items must carry who created them and who they are assigned to. Notifications must @-tag the right participant — but never the operator (who drives the system) and never the system agent.”

MANDATORY for every consumer that ships a messenger/notification surface (Herald + flavor binaries = reference impl; others inherit per §11.4.35). Detailed spec: Herald docs/design/PARTICIPANT_ATTRIBUTION.md — restated, not redefined. (A) Participant identity — every messenger MUST relate messages to a **Participant** (logical Subscriber/User); the SAME person MAY have a DIFFERENT username per messenger (logical subscriber: canonical messenger-neutral handle, kind ∈ {human, agent, service}; + per-channel aliases channel/channel_user_id/the @username used for tagging). Canonical handle closed set: Claude (reserved system-agent sentinel; never tagged) OR a subscriber @username. (B) Operator = env var, not a DB flag — the one human who drives via the agent CLI, designated by HERALD_<CHANNEL>_OPERATOR_USERNAME; a normal Participant whose handle equals that value. (C) Workable items MUST carry created_by + assigned_to (canonical handles): CLI-prompt→Operator; system/agent-detected→Claude; received-message→sender’s resolved @username; assigned_to defaults to Operator, overridable; legacy items carry "" and MUST

still parse + validate. (D) Tagging matrix: tag assigned_to if human ≠ Operator; tag created_by if human ≠ Operator ≠ Claude; NEVER tag Claude or the Operator; de-dup; resolve to the channel @username, skip if not on that channel. (E) Anti-bluff (composes §11.4): real SQLite round-trip with the new columns byte-identical (incl. legacy fixtures WITHOUT the fields); tagging matrix proven by a truth-table + a cell-flip mutation forcing FAIL; E2E real-event → real-message asserting the exact @usernames + a NEGATIVE case proving the Operator is NOT tagged; evidence under docs/qa/<run-id>/.

Composes with §11.4 + §11.4.1..§11.4.16 / §11.4.5 / §11.4.69 / §11.4.50 / §11.4.91 / §11.4.93 / §11.4.95 / §1.1. Classification: universal (§11.4.17); projects with no messenger surface inherit it latently (binds the moment they ship one, §11.4.96 pattern). Propagation gate CM-COVENANT-114-104-PROPAGATION (literal 11.4.104 across consumer fleet) + paired §1.1 meta-test mutation (strip literal → gate FAILs; gate-code = separate work item). No escape hatch — no --skip-attribution, --no-participant-tagging, --tag-operator-anyway, --attribution-later flag.

Canonical authority: constitution submodule Constitution.md §11.4.104; detailed spec Herald docs/design/PARTICIPANT_ATTRIBUTION.md. Non-compliance is a release blocker.

§11.4.105 — Natural-language intent recognition & clarification (User mandate, 2026-05-31)

Forensic anchor — verbatim user mandate (2026-05-31): “Users must NOT need to know command syntax (no COMMAND: ... prefix). They send a clear natural-language message; the System determines the intent. The System recognizes the commands it has; if none matches it infers the exact intent; if it is totally unable it replies, tags the user (@user ...), and asks to clarify precisely. We MUST always do our best to determine exact intent so we never annoy end users. This is a CORE part of the System.”

MANDATORY for every consumer that ships a messenger/command surface (Herald + flavor binaries = reference impl; others inherit per §11.4.35). Detailed spec: Herald docs/design/INTENT_RECOGNITION.md — restated, not redefined. (A) No required command syntax — users MUST NOT need to know any syntax (no COMMAND: prefix); they send plain natural language and the System determines the intent. (B) Three-tier resolution, first that succeeds wins: TIER 1 recognize the System’s existing command set from natural language → action (confident deterministic match); TIER 2 when no command matches, infer the exact intent (LLM maps language → action), NEVER guessing; TIER 3 when neither a command nor a confident intent can be determined, REPLY to the message, TAG the sender (@username, via the §11.4.104 IdentityResolver) and ask a PRECISE clarifying question NAMING the candidate intents — no guessing, no silent drop. (C) Never guess, never drop — a wrong action is worse than a clarifying question (composes §11.4.6); a message is NEVER silently dropped; only genuine ambiguity reaches Tier 3, which always replies-tags-and-asks. (D) Anti-bluff (composes §11.4): every tier ships unit + integration + E2E + full-automation tests with real captured evidence — Tier 1 truth-table (natural-language → action+fields + conservative negatives that MUST fall through to “no match”); Tier 3 E2E whose recorded reply body is EXACTLY @<sender> <specific question> + a

NEGATIVE proving a clear command does NOT trigger clarify; a paired §1.1 mutation breaking the confidence guard (false-match) OR dropping the clarify tag MUST FAIL a test; evidence under docs/qa/<run-id>/.

Composes with §11.4 + §11.4.1..§11.4.16 / §11.4.6 / §11.4.104 / §11.4.5 / §11.4.69 / §11.4.98 / §1.1. Classification: universal (§11.4.17); projects with no messenger surface inherit it latently (binds the moment they ship one, §11.4.96 pattern). Propagation gate CM-COVENANT-114-105-PROPAGATION (literal 11.4.105 across consumer fleet) + paired §1.1 meta-test mutation (strip literal → gate FAILs; gate-code = separate work item). No escape hatch — no --require-command-syntax, --guess-intent-ok, --skip-clarify, --drop-on-ambiguous flag.

Canonical authority: constitution submodule Constitution.md §11.4.105; detailed spec Herald docs/design/INTENT_RECOGNITION.md. Non-compliance is a release blocker.

§11.4.106 — Docs Chain — mechanical documentation/DB sync engine (Operator mandate, 2026-05-31)

Forensic anchor — operator mandate (2026-05-31): Docs Chain is the canonical mechanical enforcer of the documentation-sync mandates; consumers MUST use the engine instead of ad-hoc per-project doc-sync scripts, register their chains via per-context YAML, and never accept a faked transform.

MANDATORY for every consumer. Docs Chain (the vasic-digital/docs_chain engine — a universal Go bidirectional document-and-database dependency-propagation engine) is the canonical mechanical enforcer of the documentation-sync mandates. Detailed spec: the engine docs ~/Projects/docs_chain → docs/CONSTITUTION_INTEGRATION.md (distribution + inheritance + anchor mapping table) + docs/USE_CASE_CATALOGUE.md (chain recipes) — restated, not redefined. (A) Use the engine, never ad-hoc scripts — consumed by reference (the flat-layout sibling ~/Projects/docs_chain / the constitution-exposed path), inherited like §11.4.80's codegraph_* scripts: referenced, NEVER copied; ad-hoc sync_*/generate_*/summary/update_readme_doc_links scripts are superseded and retired per registered context. (B) Consumer-owned contexts — the engine is project-agnostic; the consumer registers its chains as data via .docs_chain/contexts/*.yaml (§11.4.28 decoupling); state.json + *.docs_chain.tmp are gitignored. (C) Anchors it mechanizes (per the CONSTITUTION_INTEGRATION mapping table): §11.4.12 / §11.4.53 / §11.4.45 / §11.4.56 / §11.4.57 / §11.4.59 / §11.4.60 / §11.4.65 / §11.4.86 / §11.4.93 / §11.4.95 / §12.10 / §11.4.44 — with content-hash change detection (NOT mtime, §11.4.86), atomic-rename + SQLite-txn commit + rollback (§9.2), both-dirty sync → conflict-not-silent-merge (§11.4.6), verify as the deterministic CI/pre-build gate (§11.4.50), per-run captured evidence to qa-results/docs_chain/<run-id>/ (§11.4.69). (D) NOT a replacement for authoring discipline — the source author still writes the §11.4.44 revision header; the engine only keeps exports in sync. (E) Anti-bluff (composes §11.4): the engine NEVER fakes a transform — a missing pandoc/weasyprint surfaces a typed ToolAbsentError + honest §11.4.3 SKIP-with-reason, never a fake PASS / partial write; every sync/verify carries real captured evidence.

Composes with §11.4 + §11.4.1..§11.4.16 / §11.4.6 / §11.4.28 / §11.4.80 / §9.2 / §11.4.50 / §11.4.69 / §11.4.5 / §1.1, plus the sync anchors it mechanizes (§11.4.12/.53/.45/.56/.57/.59/.60/.65/.86/.93/.95/.44, §12.10). Classification: universal (§11.4.17); projects with no derived-export/DB-sync surface inherit it

latently (binds the moment they ship one, §11.4.96 pattern). Status (§11.4.6): engine Phases 1–3 AND the CLI/YAML loader (Phase 4) IMPLEMENTED+tested — `cmd/docs_chain/main.go` registers `sync/verify/diff(alias)/doctor, go test -count=1 ./...` GREEN 7/7 packages; this CLI/loader is in the working tree but UNTRACKED and the engine is NOT yet a registered git submodule (in-tree at `./docs_chain`, HEAD = parent HEAD); only submodule distribution (Phase 6) remains PLANNED + OPERATOR-GATED — wire as phases land, claim no unshipped behaviour. Propagation gate CM-COVENANT-114-106-PROPAGATION (literal 11.4.106 across consumer fleet) + paired §1.1 meta-test mutation (strip literal → gate FAILs; gate-code = separate work item). No escape hatch — no `--ad-hoc-sync-ok`, `--skip-docs-chain`, `--fake-transform`, `--sync-evidence` optional flag.

Canonical authority: constitution submodule `Constitution.md` §11.4.106; detailed spec the Docs Chain engine `docs/docs/CONSTITUTION_INTEGRATION.md` + `docs/USE_CASE_CATALOGUE.md` (vasic-digital/docs_chain). Non-compliance is a release blocker.

§11.4.107 — Anti-bluff AV/test-validation techniques mandate (User-driven research, 2026-06-02)

Forensic (genericised, 2026-06-02): a test PASSEd on a SINGLE captured frame showing “a picture” — but it was a FROZEN/STALE frame from previously-played content (stale-producer/stuck-decoder), feature broken for the user while green; a sibling FLASHED media on the WRONG output for ~1 s before routing with no test sampling that window; a third shipped a comparator that PASSEd its own deliberately-degraded fixture (the analyzer, not the feature, was the bluff). §11.4.5 mandates captured evidence + a presence pass; §11.4.107 raises it to liveness + correct-routing + self-validated-analyzer.

Every test asserting audio/video output is genuinely playing/advancing MUST satisfy ALL of: (1) a single captured frame is NOT proof — prove LIVE, ADVANCING frames over a window via a freeze-detection oracle (near-duplicate/freeze-detect-class OR perceptual-hash adjacent-frame distance, NOT byte-identical — byte-identity is only a pre-filter); (2) an independent frame-advance counter from the platform’s compositor/decoder telemetry must increase (a different-domain oracle — flat ⇒ stuck decoder ⇒ FAIL even if pixels appear to move); (3) loading/buffering is a distinct state — wait for genuine playback-start, never false-FAIL a still-loading stream nor false-PASS a spinner (timeout+unreachable ⇒ SKIP per §11.4.3, timeout+reachable ⇒ FAIL); (4) not-stale-from-previous cross-check (new first frame ≠ previous last frame); (5) measured FPS / no-lost-frames within tolerance; (6) no-flash-on-the-wrong-output — high-frequency sampling of the non-target output during routing transitions, any content frame there ⇒ FAIL (content-protection regime classified, never guessed); (7) drive through the realistic feed/UI path, not deep-link shortcuts (UI-not-introspectable ⇒ operator-attended per §11.4.52, never fake-PASS); (8) metamorphic relations solve the oracle problem when there is no golden source (same content output-A vs output-B must match; paused ⇒ counter stops; 2× speed ⇒ ~2× advance); (9) full-reference quality metrics (SSIM/VMAF/ΔE2000) vs a golden source for owned content; (10) mutation-test every analyzer with a golden-good + golden-bad fixture pair so the analyzer itself cannot bluff (PASS on the golden-bad fixture = bluff gate — §1.1 applied to analyzers); (11) per-channel audio RMS/loudness (EBU R128) + XRUN/underrun census (a single aggregate RMS misses a dead channel); (12) OCR overlay/subtitle detection needs a per-word

confidence floor + ROI (avoids both false-FAIL and false-PASS); (13) thresholds calibrated on the project's own fixtures, not hardcoded from literature (§11.4.6). Honest gap: true photon FPS at the sink has no clean software oracle — flag it, never claim it.

4-layer per §11.4.4(b): pre-build gate (CM-AV-LIVENESS-NO-FROZEN-FRAME-class — every output-is-playing test references the liveness battery not a single frame + analyzer-self-validation gate wiring golden-good/golden-bad fixtures into meta-test) + runtime/on-device test + paired §1.1 mutation (single-frame-only → gate FAILs; analyzer PASSing its golden-bad fixture → self-validation FAILs) + HelixQA Challenge. Every PASS via `ab_pass_with_evidence` citing motion / not-stale / frame-advance / fps / per-channel-loudness / metamorphic / self-validation artefacts (`video_display/audio_output/subtitle_render` per §11.4.69) — never a single screenshot. Classification: universal (§11.4.17) — platform-neutral, reusable by ANY project validating media playback or any pixel/audio output; the project supplies its capture mechanism + calibrated thresholds per §11.4.35. Composes §11.4.5/.6/.50/.68/.69/.85 + §11.4.2/.3/.13/.48/.52/§1.1. Propagation gate CM-COVENANT-114-107-PROPAGATION (literal 11.4.107) + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.107. Non-compliance is a release blocker. No escape hatch — no `--single-frame-proves-playback`, `--skip-liveness`, `--byte-identical-freeze-OK`, `--no-frame-advance-counter`, `--allow-wrong-output-flash`, `--deep-link-shortcut-OK`, `--unvalidated-analyzer-OK`, `--aggregate-rms-suffices`, `--hardcoded-thresholds-OK` flag.

§11.4.108 — Four-layer fix-verification + runtime-signature-as-definition-of-done mandate (systematic-debugging Phase 4.5, 2026-06-03)

Forensic anchor (genericised, 2026-06-03): across one batch multiple fixes were “green” at every gate yet NOT working for the end user — a source-committed change passed by both the source pre-build gate AND the post-build gate never reached the boot-time command-line embedded in the deployed boot artifact (feature dead despite all-green); sibling fixes correctly built into the system image were masked by stale per-user overlay copies from a prior deployment (running code = stale shadow, not the fresh build); deployment did not wipe the mutable overlay so the shadow survived; validation ran against the stale shadow and reported green on code never exercised. Per `superpowers:systematic-debugging Phase 4.5`: each fix revealing a fresh “fixed-but-not-working” in a DIFFERENT place is ONE architectural VERIFICATION flaw, not independent bugs — patching symptoms is thrashing. A fix crosses FOUR distinct layers, and “fixed” at one does NOT imply fixed at the next: (1) SOURCE (committed — `grep-the-source` gate), (2) ARTIFACT (the change's BYTES landed in the produced artifact — image / bundle / installer / embedded command-line), (3) RUNTIME-ON-CLEAN-TARGET (active on a CLEAN/fresh deployment, the layer the end user experiences, no stale overlay shadowing it), (4) USER-VISIBLE (works for the end user — §11.4.5/§11.4.69 captured-evidence). The mandate (ALL hold): (1) **runtime-signature-as-definition-of-done** — a fix is DONE only when its declared runtime signature verifies on a CLEAN deployment; source-committed ≠ artifact-contains-it ≠ active-on-clean-target ≠ user-visible-working; (2) every fix declares ONE machine-checkable runtime signature (running-system property / sink-side report per §11.4.13 / §11.4.5/§11.4.69 captured-evidence / counter delta — NEVER a re-

grep of the source); this registry of per-fix runtime signatures is the SINGLE SOURCE OF TRUTH for “fixed”, REPLACING “a gate greps the source”; (3) gates span all four layers — source (pre-build), artifact (post-build — assert the BYTES landed, not merely that the source still contains the change), runtime-on-clean-target (post-deploy — assert the runtime signature), user-visible (§11.4.5/§11.4.69); (4) eliminate the stale-deployment / shadow layer by construction — deployment MUST yield a CLEAN state (wipe the mutable overlay) OR a pre-validation assertion MUST prove running-artifact == built-artifact before any validation; validation against possibly-stale state is INVALID and its PASS is a §11.4 PASS-bluff; (5) meta-rule — ≥ 3 “fixed-but-not-working” discoveries in one cycle signal an architectural VERIFICATION flaw, not three independent bugs; on the 3rd STOP patching symptoms (§11.4.4), fix the VERIFICATION pipeline, re-certify EVERY item through it on a clean target; (6) a batch is “validated” only after COMPREHENSIVE per-item runtime-signature verification on a clean baseline, NOT after the touched items’ own tests pass. Classification: universal (§11.4.17) — platform-neutral verification-pipeline disciplines reusable by ANY project that builds → deploys → validates an artifact; the project supplies its artifact format, clean-deployment / equal-artifact mechanism, and per-fix runtime-signature observables per §11.4.35. Composes §11.4.1/.2/.4/.5/.6/.27/.40/.46/.50/.52/.69/.102. Propagation gate CM-COVENANT-114-108-PROPAGATION (literal 11.4.108) + recommended per-fix gate CM-RUNTIME-SIGNATURE-REGISTRY + paired §1.1 mutations (gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md §11.4.108. Non-compliance is a release blocker. No escape hatch — no --source-green-is-done, --skip-artifact-byte-check, --validate-against-running-state, --no-clean-deployment, --skip-runtime-signature, --spot-validate-touched-only flag.

§11.4.109 — Mandatory Anti-Forgetting Enforcement: PreToolUse Guard Hook + Subagent Constitutional Preamble + Orchestrator Pre-Action Checklist (Operator mandate)

UNCONFIRMED forensic anchor (pending operator’s verbatim mandate quote). Background: emulator subagents ran raw host-direct emulator/adb because the orchestrator forgot to inject the Containers-submodule rule at dispatch time. A rule forgotten at dispatch is not enforcement. Two-layer fix: (A) constitution/scripts/hooks/guard-forbidden-commands.sh — a portable bash/awk PreToolUse hook (no jq dependency, no project-specific paths) that blocks host-direct emulator / force-push bypass / sudo / host-power at the tool-call boundary regardless of agent memory (hook = the floor); (B) constitution/docs/AGENT_GUARDRAILS.md — canonical preamble the orchestrator pastes verbatim into every subagent dispatch + Orchestrator Pre-Action Checklist (preamble = the ceiling). Consuming projects wire the hook in .claude/settings.json (or equivalent) and maintain a docs/AGENT_GUARDRAILS.md containing the SUBAGENT CONSTITUTIONAL PREAMBLE, ORCHESTRATOR PRE-ACTION CHECKLIST, and anchor literal 11.4.109. The hook is inherited by reference — NEVER copied locally. Gates: CM-ANTI-FORGETTING-ENFORCEMENT (hook + doc + test) + CM-COVENANT-114-109-PROPAGATION (literal 11.4.109 across consumer fleet). Paired §1.1 mutations (gate-code = separate work item). Classification: universal (§11.4.17). Composes §11.4.6/.10/.75/.76/.78/.79/.80/.81/.84/.98/.102/§12.

Canonical authority: Constitution.md §11.4.109; ref impl constitution/scripts/hooks/guard-forbidden-commands.sh + constitution/docs/AGENT_GUARDRAILS.md. Release blocker. No --skip-pretooluse-hook / --no-guardrails-doc / --anti-forgetting-optional / --single-layer-sufficient escape.

§11.4.110 — Pre-build build-readiness verdict + change-impact clash detection mandate (operator mandate, 2026-06-03)

Forensic anchor (genericised, 2026-06-03): a fix shipped a new system-property *read* but introduced no matching security-policy grant + no property-context type entry — the read was silently denied at runtime + the feature was dead, while every pre-build gate stayed green because the gate only grepped the source file. Generalises: a change introduces a new dependency on a *second* artifact (security policy, context file, service registry, interface freeze-snapshot, symbol table, build-graph node) that the pre-build never cross-checks. This is §11.4.108's SOURCE→ARTIFACT gap shifted LEFT to pre-build — most such clashes are statically catchable from the change diff itself. The mandate (ALL hold): (1) a single deterministic **READY-FOR-BUILD verdict** gates the rebuild (orchestrator refuses to start unless READY); (2) a **diff-driven change-impact + clash detector** cross-checks every newly-introduced second-artifact dependency (new property read ⇔ property-context type + read-grant; new service ⇔ service-context; new init service ⇔ security label; new/changed stable interface ⇔ freeze-snapshot updated; new native-lib dep ⇔ module/prebuilt resolves; new policy rule ⇔ every type/attribute defined; two batch changes on the same seam ⇔ collision acknowledged); (3) **coverage-completeness is a gate** — every changed file maps to ≥1 gate + ≥1 deployed-target test + ≥1 paired §1.1 mutation, baseline ratchets per §11.4.50; (4) **two-speed honesty** — grep-speed always-on gates vs REQUIRES_BUILD heavy gates (build-graph parse-only dry-run, full neverallow compilation, ABI diff) as diff-gated opt-in stages bounded per §12.6/§12.7; (5) every gate + wired analyzer is anti-bluff by paired §1.1 mutation; (6) **honest boundary** — a READY verdict proves static internal-consistency + ready-to-build, NOT that the feature works on the deployed target; the regime empties the *preventable* class, not the *all-defects* class (runtime/USER-VISIBLE remains §11.4.108's job). Classification: universal (§11.4.17) — platform-neutral pre-build-rigor disciplines reusable by ANY artifact-building project; the project supplies its property/service/policy registries, build-graph parse-only command, interface-freeze mechanism + changed-file→gate mapping per §11.4.35. Composes §11.4.1/4/6/9/27/50/67/75/92/108. Propagation gate CM-COVENANT-114-110-PROPAGATION (literal 11.4.110) + recommended CM-READY-FOR-BUILD-VERDICT / CM-CHANGE-IMPACT-CLASH-DETECTOR / CM-COVERAGE-COMPLETENESS-GATE / CM-BUILDGRAPH-DRYRUN-WIRED / CM-SEPOLICY-NEVERALLOW-WIRED + paired §1.1 mutations (gate-code = separate work item).

Canonical authority: Constitution.md §11.4.110. Release blocker. No --source-green-is-ready / --skip-clash-detector / --skip-coverage-gate / --no-ready-verdict / --grep-proves-neverallow / --skip-buildgraph-dryrun / --build-without-ready-verdict escape.

§11.4.111 — Resolve-by-stable-name-not-by-enumeration-index mandate (research-derived, 2026-06-03)

Forensic anchor (genericised, 2026-06-03): a platform bound an audio output to a kernel-enumerated index (card=0); a second device of a different class enumerated FIRST at boot + took slot 0, shifting the intended device to slot 1 — the static index binding now pointed at the wrong device (policy mis-attached the sink, mis-assigned its TYPE, the output switcher labelled the AV receiver “Wired Headphone” + routing collapsed to stereo). The lower layer of the SAME stack already resolved by **name** (controller-name registry) — proving the brittleness was the *index*, not the resolution capability. Generalises to every enumerated resource whose ordinal is assigned at discovery/boot/hotplug (non-deterministic across reboots, additions, topology changes). The mandate: any binding to a hardware device / resource handle / enumerated entity (audio cards, display connectors, network interfaces, storage devices, GPU render nodes, input/camera devices, container/process slots) MUST resolve by a **stable identifier** (name / UUID / serial / label / controller-name / content-hash / sink-reported identity), NOT by enumeration index / ordinal / slot, UNLESS the platform documents the ordinal as deterministically pinned AND the pin is captured + asserted in the binding. Where a stable identifier exists at one layer, every layer binding the same resource MUST use it (mixed by-name-here / by-index-there is the forbidden weak link). Honest boundary (§11.4.6): when only an ordinal exists, pin deterministically via the platform’s mechanism, capture the pin in the §11.4.108 runtime signature, document residual fragility as UNCONFIRMED: -class — never silently trust an unpinned ordinal. Classification: universal (§11.4.17) — platform-neutral binding-robustness discipline; the project supplies its stable-identifier mechanism (ALSA card name, DRM connector name, NIC predictable name, block-device UUID, etc.) + the layers that must agree per §11.4.35. Composes §11.4.6/.8/.69/.108/.110. Propagation gate CM-COVENANT-114-111-PROPAGATION (literal 11.4.111) + recommended CM-RESOLVE-BY-NAME-NOT-INDEX + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: Constitution.md §11.4.111. Release blocker. No --allow-index-binding / --ordinal-is-stable-enough / --skip-name-resolution / --trust-unpinned-index escape.

§11.4.112 — Structural-impossibility won’t-fix classification mandate (research-derived, 2026-06-03)

Forensic anchor (genericised, 2026-06-03): deep research (§11.4.8) into relocating protected (content-protection / secure-surface) video to a secondary display PROVED — from authoritative platform/HDCP docs + reproducible captured behaviour (a secure surface is blanked on any output lacking the secure flag; mirror/screencap of a secure layer returns black) — that the goal is **structurally impossible by platform design**, not a missing feature. Without a durable classification it is re-investigated every cycle (re-read the same sources, re-run the same probes, re-derive the same impossibility — compounding wasted effort). The mandate: when deep research per §11.4.8 PROVES (cited authoritative sources AND, where applicable, reproducible captured evidence) a goal is structurally impossible on the target platform (forbidden by platform design / hardware-protocol constraint / documented kernel-or-API limitation — NOT merely unimplemented or hard), the goal MUST be: (1) classified Won't-fix + closed per §11.4.90 with reason structurally-impossible; (2) documented with the impossibility evidence — cited source URLs

(per §11.4.99) + reproducible probe/captured evidence — in the tracker entry + relevant docs/ guide; (3) NOT re-attempted in future cycles — a reopen MUST cite NEW evidence the constraint changed (per §11.4.34 + §11.4.7), never re-derive the same impossibility; (4) paired with the correct posture — the entry states what the project DOES instead, so “impossible” is never confused with “broken/unhandled”. Honest boundary (§11.4.6): reserved for *proven* impossibility — “could not find a way” / “very hard” / “no time” are Operator-blocked (§11.4.21) or open work, NOT won’t-fix; mislabelling them is a §11.4 planning-layer bluff. A future platform change can make the impossible possible (durable, not eternal). Classification: universal (§11.4.17) — platform-neutral effort-conservation + honesty discipline; the project supplies the impossible goal, its constraint + cited evidence per §11.4.35. Composes §11.4.6/.7/.8/.34/.90/.99. Propagation gate CM-COVENANT-114-112-PROPAGATION (literal 11.4.112) + recommended CM-WONT-FIX-STRUCTURAL-IMPOSSIBILITY + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: Constitution.md §11.4.112. Release blocker. No --wont-fix-without-proof / --reattempt-closed-impossible / --skip-impossibility-evidence / --impossible-equals-broken escape.

§11.4.113 — Absolute no-force-push + merge-onto-latest-main mandate (User mandate, 2026-06-03)

Forensic anchor — verbatim user mandate (2026-06-03): “Any force-push is strictly forbidden! We must for every Submodule take as a base latest commit on Submodule’s main (or master) branch, then on top of it carefully to merge all changes that have to be pushed! Once all merging is carefully done we perform commit and push to all Submodule’s upstreams!” Force-push is STRICTLY FORBIDDEN with NO exception — `git push --force, --force-with-lease, +<ref>`, any history-rewriting overwrite of a remote ref — against EVERY repository this Constitution governs (main repo, this constitution submodule, every owned + nested submodule, every upstream). No operator-approval path, no “after merge-first audit” path. Mandated 6-step integration for any repo/submodule with commits to publish OR diverged mirrors: (1) `git fetch --all --prune --tags` all remotes; (2) base = LATEST commit on canonical main/master (most-advanced mirror tip); (3) carefully MERGE every change on top — union, preserve BOTH sides, NEVER `-s ours` / rebase / reset that drops commits (§9 no-commit-loss); (4) resolve conflicts carefully — no markers, no file dropped, gates/tests still pass; (5) commit the merge (stage only intended files, NEVER `git add -A` in a submodule per §11.4.30); (6) push to ALL upstreams — each is a fast-forward because the merge commit descends from every mirror tip, so NO force is needed (an upstream still rejecting → return to step 1 for it, merge its new tip, re-validate, re-push). TIGHTENS §11.4.41 / §11.4.71 / §9.2 / CONST-043 — force-push escape hatch REMOVED: even WITH operator approval, even after a clean merge-first audit, force-push is forbidden, because the merge-onto-latest-main path is always available so force is never necessary; those clauses’ merge-first/fetch-first machinery stays, only their terminal “...then force-push” step is struck. Classification: universal (§11.4.17). Composes §2.1 / §9 / §9.2 / §11.4.4 / §11.4.6 / §11.4.26 / §11.4.37 / §11.4.40 / §11.4.41 / §11.4.71 / §11.4.88 / CONST-043. Propagation gate CM-COVENANT-114-113-PROPAGATION (literal 11.4.113) + recommended CM-NO-FORCE-PUSH-ABSOLUTE (scan tracked scripts/hooks for push `--force` / `--force-with-lease` / push `+<ref>` + reject; §11.4.109-class PreToolUse guard blocks the class) + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: Constitution.md §11.4.113. Release blocker. No `--force / --force-with-lease / --allow-force-push / --force-push-authorised / --skip-merge-onto-main` escape.

§11.4.114 — Last-known-good-tag regression isolation mandate (1.1.8-dev remediation, 2026-06-03)

When a previously-working feature/behaviour is observed broken, the FIRST diagnostic action MUST be to identify the last release tag (or commit/build/deploy) at which it was KNOWN-GOOD and diff/bisect the broken state against it — BEFORE any open-ended root-cause hunt or speculative fix. The known-good revision is the regression oracle: bounds the search to commits between good and now (`git diff <good-tag>..HEAD --stat` of the feature's files = captured evidence), marks it a regression REPAIR not from-scratch design, gives each suspect file a behavioural oracle. An operator-volunteered known-good tag is load-bearing → act on it first. Default to SURGICAL forward-fix (keep post-good-tag features, revert ONLY the broken sub-part) over wholesale revert (which loses the batch's other working features) unless operator prefers wholesale. Honest boundary (§11.4.6): “it worked before” is a HYPOTHESIS until the tag is identified AND the feature confirmed working there; “probably regressed last batch” without the diff is a guess; a never-worked feature is not a regression. Classification: universal (§11.4.17). Composes §11.4.4 / §11.4.6 / §11.4.7 / §11.4.40 / §11.4.43 / §11.4.102 / §11.4.108. Propagation gate CM-COVENANT-114-114-PROPAGATION (literal 11.4.114) + recommended CM-REGRESSION-ISOLATED-AGAINST-KNOWN-GOOD + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: Constitution.md §11.4.114. Release blocker. No `--skip-known-good-diff / --root-cause-from-scratch / --assume-regression-source / --wholesale-revert-without-isolation` escape.

§11.4.115 — RED-baseline-on-the-broken-artifact + polarity-switch mandate (1.1.8-dev remediation, 2026-06-03)

Strict refinement of §11.4.43's RED step. Every RED test MUST REPRODUCE the defect on the CURRENT pre-fix artifact (the actual broken build/deploy), capturing positive evidence per §11.4.5 / §11.4.69 / §11.4.107 that the defect is genuinely present — never a synthetic failure the fix is then written to agree with. The SAME source MUST carry a single polarity switch (env flag / parameter, canonical RED_MODE, default 1 = reproduce-and-assert-defect-present) that flipped to 0 post-fix becomes the GREEN regression-guard asserting the defect is ABSENT. One source, two roles: the bug-catcher IS the regression-guard; no separate happy-path test is the primary guard (that only proves the test agrees with the fix — the §11.4.43 PASS-bluff). RED-on-broken then GREEN-on-fixed (clean target per §11.4.108) both captured. Honest boundary (§11.4.6): a RED run that does NOT fail on the broken artifact is a finding (close per §11.4.7 with negative evidence, or fix the blind test). Classification: universal (§11.4.17). Composes §11.4.1 / §11.4.2 / §11.4.5 / §11.4.69 / §11.4.107 / §11.4.4 / §11.4.7 / §11.4.43 / §11.4.50 / §11.4.108 / §11.4.114. Propagation gate CM-COVENANT-114-115-PROPAGATION (literal 11.4.115) + recommended CM-RED-POLARITY-SWITCH-PRESENT + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: Constitution.md §11.4.115. Release blocker. No --green-only-test / --skip-red-reproduction / --separate-happy-path-suffices / --synthetic-red-OK escape.

§11.4.116 — Real-time conductor[?]autonomous-test-framework sync channel mandate (1.1.8-dev remediation, 2026-06-03)

Any autonomous, long-running test/QA/validation framework an external orchestrator (conductor agent / operator) depends on for real-time decisions MUST expose a real-time sync channel: (1) a structured append-only event stream (JSONL or equivalent, one event per line, never rewritten — session-start / phase-transition / per-test-or-challenge-start / captured-evidence-path / external-call (LLM/vision/sink-probe) / error / per-item-verdict); (2) an atomically-rewritten status snapshot (single small file, write-temp-then-rename so no torn read) with current session/phase/item + counters + last verdict. Verdicts use the closed PASS / FAIL / SKIP / OPERATOR-BLOCKED vocabulary (§11.4.45). The conductor tails it live → real-time sync, §11.4.4 interrupt on a fresh defect, never idles blindly (§11.4.94 / §11.4.97). Anti-bluff: every verdict event MUST carry the evidence path that backs it (§11.4.69) — a PASS event with no evidence path is a channel-layer bluff; a snapshot PASS contradicted by a stream with no evidence event → treat as FAIL; no start-event → no verdict-event. Owned project-agnostic submodule (§11.4.28): channel stays project-neutral, the consumer registers its data via the public API at runtime, never hardcoded. Classification: universal (§11.4.17). Composes §11.4.4 / §11.4.5 / §11.4.69 / §11.4.27 / §11.4.28 / §11.4.45 / §11.4.52 / §11.4.89 / §11.4.94 / §11.4.97. Propagation gate CM-COVENANT-114-116-PROPAGATION (literal 11.4.116) + recommended CM-AUTONOMOUS-FRAMEWORK-SYNC-CHANNEL + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: Constitution.md §11.4.116. Release blocker. No --no-sync-channel / --end-of-session-report-suffices / --verdict-without-evidence-path / --torn-status-write-OK escape.

§11.4.117 — Computer-vision / OCR pixel-oracle fallback for non-introspectable UIs mandate (1.1.8-dev remediation, 2026-06-03)

Any test that drives a UI control OR asserts on-screen content MUST NOT assume the accessibility/semantic/DOM hierarchy is the source of truth for what the user sees. When the hierarchy is blank/partial/known-unreliable (TV-Compose, leanback, canvas/SurfaceView/GL, games, custom-rendered UIs), the test MUST fall back to a PIXEL ORACLE: (1) DRIVE input by computer-vision template-match (locate a control by rendered appearance → tap coordinates), not a node id; (2) ASSERT content by ROI OCR (read the rendered text the user reads) with a per-word confidence floor + ROI per §11.4.107(12), not a hierarchy text attribute that may be absent/stale. The tool MUST both drive input AND read pixels — a hierarchy-only tool is NOT a content oracle. Makes §11.4.52's "near-empty hierarchy → INFEASIBLE" constructive: pixel-drive, not only SKIP. Anti-bluff (§11.4.107(10)): the CV/OCR analyzer is self-validated (golden-good PASS, golden-bad FAIL, wired into meta-test); thresholds calibrated on the project's own frames, not hardcoded (§11.4.6). Honest boundary: BOTH hierarchy blank AND pixel oracle infeasible (secure surface black-captures §11.4.112, geo-unreachable) → SKIP-with-reason §11.4.3 + tracked operator-attended migration §11.4.52, never fake PASS. Classification: universal (§11.4.17). Composes §11.4.5 / §11.4.6 / §11.4.48 / §11.4.49 / §11.4.52 / §11.4.107 / §11.4.112. Propagation gate CM-

COVENANT-114-117-PROPAGATION (literal 11.4.117) + recommended CM-CV-OCR-PIXEL-ORACLE-FALLBACK + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: Constitution.md §11.4.117. Release blocker. No --hierarchy-is-content-oracle / --skip-pixel-fallback / --unvalidated-ocr-OK / --hardcoded-ocr-threshold-OK escape.

§11.4.118 — Discovery-pressure to confirm known-issue-set completeness mandate (1.1.8-dev remediation, 2026-06-03)

A remediation/release cycle MUST NOT treat “every reported defect is fixed” as “the build is good” — the reported set is biased by what the operator happened to test (“we only see what we test”). After/alongside fixing the reported set, run a discovery + stress pass across ALL target devices/environments that deliberately exercises subsystems/journeys/edge-cases BEYOND the reported defects — to CONFIRM the reported set is the COMPLETE critical set + surface unreported defects before the user does. The pass MUST produce PROVABLE coverage: an enumerated list of subsystems/journeys/stress-scenarios exercised, each with outcome (no-new-issue, or a new tracker entry per §11.4.15 / §11.4.16). “We found no other issues” is a §11.4 bluff unless paired with “here is the enumerated set we exercised” — absence of evidence of looking is not evidence of absence. New findings trigger §11.4.4 interrupt + §11.4.114/§11.4.115 isolation→RED→fix. Honest boundary (§11.4.6): discovery reduces the unknown-unknown surface, does not prove zero remaining defects — earned claim is “reported set + enumerated discovery set addressed,” un-exercised subsystems stated as honest coverage gaps (§11.4.3), never silently implied clean. Classification: universal (§11.4.17). Composes §11.4.4 / §11.4.5 / §11.4.69 / §11.4.25 / §11.4.40 / §11.4.42 / §11.4.52 / §11.4.85 / §11.4.114 / §11.4.115 / §11.4.119. Propagation gate CM-COVENANT-114-118-PROPAGATION (literal 11.4.118) + recommended CM-DISCOVERY-COVERAGE-ENUMERATED + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: Constitution.md §11.4.118. Release blocker. No --reported-set-suffices / --skip-discovery-pass / --no-issues-without-coverage-list / --assume-complete escape.

§11.4.119 — Single-resource-owner partitioning for parallel hardware testing mandate (1.1.8-dev remediation, 2026-06-03)

Strict refinement of §11.4.58 / §11.4.103 for hardware contention. When multiple parallel streams exercise SHARED hardware / any exclusive-access resource (a media-playback path, a single HDMI/audio sink, an exclusive device handle, a serial/JTAG line, a GPU under exclusive capture), exactly ONE stream MUST own each such resource at a time. The owner drives it (playback, input injection, capture); every other concurrent stream targeting it MUST be READ-ONLY (passive probes — dumphys / /proc / /sys reads / sink-side network probes / log tails). Parallelism partitioned by resource: distinct devices/sinks/handles fully concurrent (stream-per-device), but the same device’s exclusive resource is single-owner. Ownership enforced by an advisory lock/token (§11.4.58 L3 + the hardware analogue of §11.4.84 quiescence), event-driven (claim when the resource frees, release the moment work completes so the next queued stream claims it per §11.4.103 auto-backfill). Why: concurrent drivers of one exclusive resource produce

CROSS-CONTAMINATED evidence (sink reports the wrong stream's audio, foreground belongs to whichever am start landed last, input events interleave) — a PASS under contention is a §11.4 evidence-integrity bluff. Honest boundary (§11.4.6): a “read-only” probe stream MUST issue NO state-changing command against a device it does not own; a non-partitionable resource → streams serialize (single-owner over time), not run concurrently and hope. Classification: universal (§11.4.17). Composes §11.4.5 / §11.4.69 / §11.4.13 / §11.4.50 / §11.4.58 / §11.4.82 / §11.4.84 / §11.4.103 / §11.4.118. Propagation gate CM-COVENANT-114-119-PROPAGATION (literal 11.4.119) + recommended CM-SINGLE-RESOURCE-OWNER-PARTITION + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: Constitution.md §11.4.119. Release blocker. No --allow-concurrent-resource-drivers / --skip-ownership-lock / --read-only-may-mutate / --contended-evidence-OK escape.

§11.4.120 — Fix-breaks-its-own-gate reconciliation mandate (1.1.8-dev remediation, 2026-06-03)

When a correct fix causes a pre-existing gate/test to FAIL because that gate asserted the OLD (now-removed-or-changed) behaviour, the gate FAIL is the CORRECT signal the fix landed — it MUST NOT be suppressed by either forbidden response: (1) FAKE-PASSING the gate (edit to always pass, weaken to a tautology, or delete — a §11.4 gate-layer bluff + a §11.4.84 mutation-residue risk); (2) REVERTING the correct fix to satisfy the stale gate (re-introducing the defect). Required response = RECONCILIATION: rewrite the gate to assert the NEW mechanism the fix introduced, backed by captured evidence of the new correct behaviour, AND update its paired §1.1 mutation so the mutation breaks the NEW invariant. Discriminator vs bluffing: after reconciliation gate + mutation still form a valid §1.1 pair (mutate the new invariant → gate FAILs); a bluffed gate's mutation no longer makes it FAIL. The reconciliation MUST be a visible evidence-cited change, never a silent assertion-weakening. Honest boundary (§11.4.6): a post-fix gate FAIL is NOT automatically “stale gate” — investigate per §11.4.102 first; the FAIL may be the gate correctly catching a REGRESSION the fix introduced (then the FIX is wrong). Reconcile ONLY when investigation PROVES the gate asserted old-correct-now-removed behaviour AND the new behaviour is intended + evidence-confirmed. Classification: universal (§11.4.17). Composes §11.4.1 / §11.4.4 / §11.4.6 / §11.4.84 / §11.4.102 / §11.4.108 / §1.1. Propagation gate CM-COVENANT-114-120-PROPAGATION (literal 11.4.120) + recommended CM-GATE-RECONCILED-NOT-FAKE-PASSED + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: Constitution.md §11.4.120. Release blocker. No --fake-pass-stale-gate / --revert-fix-for-gate / --weaken-assertion-to-pass / --delete-failing-gate escape.

§11.4.121 — No-commit-while-build-writes-tracked-artifacts mandate (1.1.8-dev remediation, 2026-06-03)

A commit (especially git add -A / any broad stage) MUST NOT run while a build/packaging/generation step is actively writing artifacts INTO tracked (version-controlled) directories — it races the writer and stages a PARTIAL or stale artifact (a §11.4.108 SOURCE→ARTIFACT integrity failure landed in version control). The commit MUST be deferred until the build step that writes tracked artifacts has COMPLETED, so the tree is quiescent at the artifact layer AND the committed

artifacts are the FRESH, whole outputs (no stale pre-rebuild artifact, no half-written file). Before committing tracked build outputs, verify the writing step finished (process exit / completion marker / per-artifact mtime $\geq$ build-start). Build-output analogue of §11.4.84 (no commit while a mutation gate is in flight): both close “commit captures transient non-final tree state” at two write-sources. Gitignored build outputs (out/ / dist/) are exempt — the race doesn’t apply. Honest boundary (§11.4.6): “the build probably finished” is not “the build finished” — verify with a completion signal; source changes NOT touching the build’s tracked-artifact dirs are fine mid-build. PROJECT-SPECIFIC instance (which tracked dir + which build steps write it) recorded in the consuming project’s own governance per §11.4.35. Classification: universal (§11.4.17). Composes §11.4.6 / §11.4.30 / §11.4.58 / §11.4.84 / §11.4.88 / §11.4.96 / §11.4.103 / §11.4.108. Propagation gate CM-COVENANT-114-121-PROPAGATION (literal 11.4.121) + recommended CM-NO-COMMIT-DURING-ARTIFACT-WRITE + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: Constitution.md §11.4.121. Release blocker. No --commit-during-build / --stage-partial-artifact-OK / --assume-build-finished / --skip-build-completion-check escape.

§11.4.122 — No-silent-removal-of-existing-components-without-operator-confirmation mandate (User mandate, 2026-06-03)

Forensic anchor — verbatim user mandate (2026-06-03): “Never ever remove any application, system component or service from already existing codebase / System without interactively asked question to us! THIS IS MANDATORY RULE / CONSTRAINT!” Case study (FACT): in the 1.1.8-dev burn-down, F2 (an Apple-TV-class app) + F4 (a Huawei HMS component) were removed WITHOUT asking; the operator reversed both.

No application / system component / service / package / feature / driver / module / library / prebuilt asset — any already-existing end-user capability — may be removed (deleted, dropped from the package set, disabled-into-non-shipping, unbundled, de-listed) from the existing codebase / shipped System WITHOUT FIRST interactively asking the operator (via the §11.4.66 interactive mechanism — AskUserQuestion on Claude Code, never free-text, never silent, never an autonomous decision) and receiving an EXPLICIT keep-or-remove decision. A removal is: a deletion from the package/service/module manifest set whose NET EFFECT is “a capability that shipped before no longer ships.” Adding / replacing-with-operator-approved-equivalent / fixing is NOT a removal; when uncertain, treat AS a removal and ask (§11.4.6 no-guessing). A silent removal is a release blocker regardless of rationale (dedup / “broken anyway” / geo-restricted / incompatible / superseded). The tracked DROP path: ask → operator approves → mark item Obsolete (→ Fixed.md) with reason feature-removed + operator-approval citation (§11.4.90) → then remove. Classification: universal (§11.4.17). Composes §11.4.66 / §11.4.101 (removal of a live capability is exactly the high-blast-radius class §11.4.101 reserves for an operator question, never an autonomous decision) / §11.4.90 / §11.4.112 / §11.4.6 / §11.4.40 / §11.4.42. Propagation gate CM-COVENANT-114-122-PROPAGATION (literal 11.4.122) + recommended gate CM-NO-SILENT-COMPONENT-REMOVAL + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: Constitution.md §11.4.122. Release blocker. No --remove-without-asking / --silent-removal / --autonomous-removal-OK / --dedup-removal-exempt / --it-was-broken-anyway escape.

§11.4.123 — Rock-solid-proof-or-deep-research mandate (User mandate, 2026-06-03)

Forensic anchor — verbatim user mandate (2026-06-03): “Every single reported issue MUST BE fully and 100% validated with rock solid proofs! Nothing can be considered fixed or completed without hard evidence! No false results or bluff(s) of any kind is allowed! If we are not sure on how to achieve full testing, validation and verification of something we MUST ALWAYS perform deep web research for all possible data (articles, documentation, guides, and other resources) and opensourced codebases which we can use to solve our problems and perform testing with validation and verification which produces rock-solid evidence(s) and leaves no space for false results or any kind of bluff!” Case study (FACT): in the 1.1.8-dev remediation the validation method for FLAG_SECURE secondary-display capture (black frames) + non-introspectable streaming-app UIs (blank accessibility tree) was unclear; deep web research (docs/research/testing_frameworks_20260603/) yielded the CV/OCR/liveness/sink-probe oracle stack (now §11.4.107/§11.4.112/§11.4.117), making rock-solid evidence possible — “unclear how to validate” is a research trigger, never a bluff licence.

Every reported issue / fix / claimed completion MUST be 100% validated with rock-solid CAPTURED proof (§11.4.5/§11.4.69/§11.4.107) before it is marked fixed/implemented/completed (§11.4.33). Nothing is fixed/complete without hard captured evidence — metadata-only / config-only / absence-of-error / grep-without-runtime PASS forbidden (§11.4/§11.4.1); no false results, no bluff of any kind. Operative addition: when UNSURE how to fully test/validate/verify something (no obvious evidence-producing method, OR the candidate yields only metadata/config/absence-of-error evidence), the agent MUST ALWAYS FIRST perform deep web research per §11.4.8+§11.4.99 (official docs, articles, guides, vendor refs, standards, issue trackers, reusable open-source codebases) to DISCOVER or BUILD an evidence-producing validation method that leaves no room for a false result. Declaring “untestable”/“not automatable” or accepting a metadata-only PASS WITHOUT first exhausting this deep-research path is itself a §11.4.123 violation (same severity as a PASS-bluff); the cited research (URLs + the method, OR “NO external solution found — original work” per §11.4.8) is the proof the path was exhausted. Only after that research genuinely fails may the item be PENDING_FORENSICS:/Operator-blocked (§11.4.21)/structurally-impossible won’t-fix (§11.4.112) with the cited research as earned evidence. Classification: universal (§11.4.17). Composes §11.4.5/§11.4.6/§11.4.8/§11.4.52/§11.4.69/§11.4.99/§11.4.107/§11.4.118/§11.4.21/§11.4.112. Propagation gate CM-COVENANT-114-123-PROPAGATION (literal 11.4.123) + recommended gate CM-ROCK-SOLID-PROOF-OR-RESEARCH + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: Constitution.md §11.4.123. Release blocker. No --metadata-pass-suffices / --skip-proof / --untestable-without-research / --config-only-closure-OK / --bluff-when-unsure escape.

§11.4.124 — Dead/unwired-code investigate-before-remove mandate (User mandate, 2026-06-04)

Forensic anchor — verbatim user mandate (2026-06-04): “Before removing any seemingly-dead (zero-importer / unwired) codebase, we MUST investigate via git history where/how it was originally used and how it became dead. Removal is permitted ONLY when we have captured PROOF it is genuinely no longer needed — and that removal MUST be its own separate commit with a proper descriptive message. If there is no such proof, the code MUST be investigated for where/how it should be wired in properly, and any missing or unwired tests MUST be added. We MUST ALWAYS be extra careful with any codebase removal.”

“Zero importers / never called / unwired $\Rightarrow$ dead $\Rightarrow$ delete” is a GUESS (§11.4.6), never a finding — a “no references” result proves only current non-reference, not genuinely-unneeded. Before removing ANY seemingly-dead element (zero-importer / never-called / unwired function / method / type / file / module / package / asset / config / build target) the agent MUST FIRST investigate via git history (`git log --follow`, `git log -S/-G` pickaxe, blame on the deleted call-site) and capture as FACT: (1) WHERE/HOW it was originally wired in, (2) WHEN/HOW it became dead (call-site deleted deliberately / by mistake-regression / never-completed / refactored-unreachable), (3) whether “no references” is real OR a hidden reference the static tool cannot see (reflection / dynamic dispatch / build-tags / codegen / DI / plugin registry / FFI / config-driven). Removal is permitted ONLY with captured PROOF the element is genuinely no longer needed; that removal MUST be its OWN SEPARATE COMMIT (independently reviewable + revertible per §11.4.84/ §11.4.92) with a descriptive message citing the git-history evidence (+ §11.4.122 operator-confirmation if it is an end-user capability; tracked via §11.4.90 Obsolete ($\rightarrow$ Fixed.md)). With NO such proof the element MUST NOT be removed — instead investigate WHERE/HOW to wire it in properly (restore a mistakenly-deleted call-site per §11.4.114; finish never-completed wiring) AND add any missing / unwired tests (§11.4.27/§11.4.43/§11.4.115). ALWAYS be extra careful with removal: when uncertain, default to NOT removing (investigate + wire + test) per §11.4.6 + §11.4.101 + §11.4.122; “probably dead” is never sufficient — the bar is captured proof. Classification: universal (§11.4.17). Composes §11.4.6 / §11.4.8 / §11.4.84 / §11.4.90 / §11.4.92 / §11.4.101 / §11.4.114 / §11.4.122 / §11.4.27 / §11.4.43 / §11.4.115. Propagation gate CM-COVENANT-114-124-PROPAGATION (literal 11.4.124) + recommended gate CM-DEAD-CODE-INVESTIGATE-BEFORE-REMOVE (a net-deletion commit must be removal-only + cite the git-history investigation OR be part of a tracked Obsolete item) + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: Constitution.md §11.4.124. Release blocker. No `--zero-importers-means-dead` / `--delete-unwired-on-sight` / `--skip-git-history-investigation` / `--remove-without-proof` / `--bundle-removal-with-other-work` escape.

§11.4.125 — Code-review-agent gate before pre-build + main build (mandatory multi-layer review) (User mandate, 2026-06-04)

Forensic anchor — verbatim user mandate (2026-06-04): “After all fixes/ changes/implementations are done, BEFORE running pre-build tests and the main build, dispatch code-review agent(s) that analyze all work done + all existing data/ facts + the existing codebase + current git history to determine quality, safety, and

whether the fixes/changes will REALLY work; they MUST validate and verify that every test covering the fixes/changes genuinely validates the work with NO chance of false results or bluff of any kind. Any finding MUST be fixed, polished, improved, and covered with additional tests before the build proceeds. Multiple strong layers of checks.”

After all fixes / changes / implementations in a batch are done, and BEFORE running the pre-build test sweep AND the main (artifact) build (for ANY project), the agent MUST dispatch one or more dedicated code-review agent(s) (subagent-driven by default per §11.4.70/§11.4.20) performing a multi-layer review that: (1) analyzes ALL work done in the batch (every fix/change + its source diff + stated intent); (2) analyzes ALL existing data + facts (captured evidence per §11.4.5/§11.4.69/§11.4.107, tracker entries, prior findings, the §11.4.108 runtime-signature registry); (3) analyzes the existing codebase (blast radius per §11.4.92, cross-feature interaction, contract integrity of every dependency); (4) analyzes current git history (what each change touched, how it composes with concurrent/recent work, whether it reproduces a known-broken pattern per §11.4.114/§11.4.124); (5) determines quality + safety + will-it-REALLY-work (robust + not error-prone — no solve-A-create-B; no host/data/security regression; genuinely delivers the end-user-visible behaviour per §11.4/§107); (6) validates + verifies the tests covering the work — every covering test genuinely exercises the work-under-test and catches its negation, with ZERO chance of a false result or bluff (a test that PASSES on broken-for-the-user work, a metadata-only/config-only/absence-of-error/grep-without-runtime assertion, or a gate whose paired §1.1 mutation does not make it FAIL is a finding). Any finding (defect / error-prone change / safety risk / will-not-really-work / bluff-or-false-result-capable test / missing-coverage gap) MUST be fixed, polished, improved, and covered with additional tests (four-layer per §11.4.4(b), TDD-RED-first per §11.4.43/§11.4.115) BEFORE the pre-build sweep + main build proceed; the review iterates (re-review after each remediation) until no blocking findings remain. The review is itself anti-bluff (its conclusions are captured evidence per §11.4.5/§11.4.69; a rubber-stamp review of a defective batch = PASS-bluff). It is one of MULTIPLE STRONG LAYERS — complementing, never replacing, the §1 pre-build sweep, §11.4.92 multi-pass (author-side self-review; §11.4.125 adds the structurally-separated reviewer seam per §11.4.70), §11.4.108 four-layer fix-verification, §11.4.110 build-readiness verdict, and the post-build / runtime-on-clean-target / user-visible layers. Composes §11.4 / §11.4.1 / §11.4.4 / §11.4.6 / §11.4.40 / §11.4.43 / §11.4.50 / §11.4.70 / §11.4.20 / §11.4.92 / §11.4.102 / §11.4.107 / §11.4.108 / §11.4.110. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-125-PROPAGATION (literal 11.4.125) + recommended gate CM-CODE-REVIEW-GATE-BEFORE-BUILD (build starts only with a fresh code-review-completed marker for the current batch, produced after the last fix + before the pre-build sweep + main build) + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: Constitution.md §11.4.125. Release blocker. No `--skip-code-review` / `--build-without-review` / `--no-review-gate` / `--review-optional` / `--trust-the-author` escape.

§11.4.126 — Default autonomous-loop working mode from first prompt (User mandate, 2026-06-04)

Forensic anchor — verbatim user mandate (2026-06-04): “Make sure that you continue work in endless fully autonomous loop, do not stop until new fully validated and verified version (tag) is created and published (all submodules and

main repo) or IN A CASE OF some other main stream work until it is fully completed with all side work streams and nothing else is left in our working queue! THIS MUST BE ALWAYS the default working mode without us asking you! We tend to achieve ABSOLUTE EFFICIENCY, with this and all other projects which will incorporate this MANDATORY RULE / CONSTRAINT!!! This way of (your) working will be ALWAYS applied / followed / executed / fully respected, as soon as we assign / send first request (prompt) in the session! This stops only if we explicitly say so or nothing is left to be done in current working scope (release that will come / upcoming version)!!! Any mimicking (imitation) of this behavior / rules / mandatory constraints, false results or any kind of bluff(s) is ABSOLUTELY FORBIDDEN!!!”

The endless fully-autonomous loop is the **DEFAULT working mode**, engaged automatically the moment the operator sends the FIRST request / prompt of a session — the operator MUST NOT have to ask for it, request it, restate it, or re-enable it per session. §11.4.87 framed the endless-loop covenant as an explicit-instruction opt-in; §11.4.126 is the **capstone** that promotes the same covenant to always-on: from the first prompt onward, every agent operates in the §11.4.87 loop discipline as the standing default, with §11.4.94 zero-idle, §11.4.97 maximum-idle-use, §11.4.101 autonomous-decision-over-blocking, and §11.4.103 continuous-parallel-stream all engaged by default — no per-session activation handshake. The continuation contract: the loop continues until ONE of two terminal conditions holds — (A) **Release scope** — a new, fully-validated-and-verified version (tag) is created AND published across all owned submodules AND the main repo to all configured remotes (per §2.1 + §11.4.40 + §11.4.113); OR (B) **Non-release main-stream scope** — the main-stream goal is fully completed AND every side work stream is done AND the working queue holds nothing left for the current scope. Until (A) or (B) holds, the agent MUST keep working per §11.4.42 / §11.4.72 / §11.4.94 / §11.4.103. The loop STOPS ONLY on: (1) the operator explicitly saying so (STOP / pause / end); (2) nothing left to do in the current working scope with the queue genuinely empty per the (A)/(B) terminal conditions; (3) a §12 host-session-safety demand. Idle-while-blocked parks one work unit, it does not stop the loop (§11.4.101 + §11.4.94 + §11.4.97). Goal — ABSOLUTE EFFICIENCY; applies to this project AND every project that incorporates this Constitution. Anti-bluff: mimicking / imitating this loop behaviour, narrating continuation without performing it, fabricating progress, or emitting false / bluff results of ANY kind is ABSOLUTELY FORBIDDEN (composes §11.4 / §11.4.1 / §11.4.2 / §11.4.5 / §11.4.6 / §11.4.50 / §11.4.69 / §11.4.107); a report claiming the loop ran while no real work / no captured evidence was produced is a §11.4 PASS-bluff at the operating-mode layer. Classification: universal (§11.4.17). Composes with §11.4.87 / §11.4.94 / §11.4.97 / §11.4.101 / §11.4.103 / §11.4.66 / §11.4.6 / §11.4.40 / §11.4.42 / §11.4.72 / §11.4.113 / §2.1 / §12. Propagation gate CM-COVENANT-114-126-PROPAGATION (literal 11.4.126) + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: Constitution.md §11.4.126. Release blocker. No --ask-before-continuing / --single-turn-only / --not-default-loop / --mimic-OK escape.

§11.4.127 — Session-handoff resumption-prompt mandate (User mandate, 2026-06-06)

Forensic anchor — verbatim user mandate (2026-06-06): “make sure that in situations like this now when new session is needed you ALWAYS prepera such sentence - which will be valid for particular moment and the phase of the project and enough for work to continue.”

When a fresh session is needed (context limits / degradation) OR the operator asks whether a new session is needed / requests a handoff, the agent **MUST ALWAYS** prepare + proactively provide a ready-to-paste **resumption prompt valid for that EXACT moment + project phase** — self-contained for ZERO-loss resume. Two variants: SHORT first-sentence (“Read <handoff docs>, then continue <goal> ...”) + FULL block on demand. **MUST** (1) point to live handoff docs (.remember/remember.md if present + docs/CONTINUATION.md §12.10, read FIRST + git fetch --all); (2) state PHASE + NEXT action + terminal goal; (3) embed exact live-state anchors (build IDs/MD5, device/target serials, commit HEAD, in-flight PIDs+log paths, captured-evidence paths); (4) restate binding constraints (anti-bluff §11.4, no-force-push §11.4.113, exact version/naming, hardware/target gotchas); (5) be MOMENT-VALID never generic. Handoff docs current BEFORE the prompt (§12.10). Missing/stale/generic = §11.4.127 violation. Composes §12.10 / §11.4.6 / §11.4.66 / §11.4.87 / §11.4.103 / §11.4.126. Classification: universal (§11.4.17). Gate CM-COVENANT-114-127-PROPAGATION (literal 11.4.127) + paired §1.1 mutation.

Canonical authority: Constitution.md §11.4.127. Release blocker. No --skip-handoff-prompt / --generic-prompt-OK / --no-resumption-sentence / --handoff-without-state escape.

§11.4.128 — Always-on device-recording mandate (User mandate, 2026-06-06)

Forensic anchor — direct user mandate (2026-06-06): ALWAYS live-record all available data from all test/manual-testing devices, EXTRA carefully (never harm device / performance / no side effects); raw NOT processed without need (token-conscious); raw ALWAYS git-ignored + code-intelligence-excluded; only curated evidence committed, only at release prep.

For EVERY test/debug device + every device under known manual testing, across EVERY transport (USB / wireless ADB / SSH / serial / network introspection API), ALWAYS live-record all analysable data: activities, all logs, perf metrics (CPU/mem/I/O/thermal/load), every sink-side report (§11.4.13), any live-changeable param. (1) **Side-effect-free** — non-invasive read-only probes, bounded sampling + write-volume, observer-effect budget; a perturbing recorder = §11.4.128 violation. (2) **Background + parallel + subagent-driven** (§11.4.103 + §11.4.70) — never blocks main stream. (3) **Token-conscious record-now-analyse-later** — raw NOT processed without need; only analyse-trigger = release-tag prep (§11.4.40/§11.4.42) OR explicit operator ask. (4) **Raw git-ignored (w/ §11.4.77 regen declaration) + code-intelligence-excluded (§11.4.78/.79)** — only curated evidence committed, only at release prep under docs/qa/<run-id>/ (§11.4.83). (5) **Layout** <recording-root>/YYYY-MM-DD/<combined main+submodules hash>/<DEVICE>_<SERIAL>/recording_NNN/<files>. (6) **Anti-bluff** — recorder-running-but-no-growing-corpus = §11.4 bluff; curated findings trace to real raw paths; recorder health = captured evidence (§11.4.5/.69). Composes §11.4.2/.5/.13/.69/.40/.42/.70/.77/.78/.79/.83/.103/.119. Classification: universal (§11.4.17). Gate CM-COVENANT-114-128-PROPAGATION (literal 11.4.128) + rec. gate CM-DEVICE-RECORDING-ALWAYS-ON + paired §1.1 mutation.

Canonical authority: Constitution.md §11.4.128. Release blocker. No --skip-recording / --record-without-layout / --commit-raw-corpus / --index-raw-corpus / --analyse-corpus-always / --invasive-probe-OK escape.

§11.4.129 — Huge-blocker release protocol (User mandate, 2026-06-06)

Forensic anchor — direct user mandate (2026-06-06): on a huge blocker during release validation: STOP all testing → fix ALL discovered issues → process all recorded data from last session → land rock-solid fixes → author NEW validation+verification tests of ALL supported test types → rebuild → reflash → RESTART full validation+verification of every fix/change from the last release tag to now — both devices in parallel, recorded, real physical proofs, no bluff.

On a HUGE BLOCKER (release-blocking-severity: core capability broken / regression invalidating the cycle / blast radius reaching the batch's other fixes), in order, NO shortcut: (1) STOP all testing every device (§11.4.4 STOP at release granularity). (2) Fix ALL discovered issues (not just the blocker) — root-cause §11.4.102 + isolate against last known-good tag §11.4.114. (3) Process all recorded data from last session (the §11.4.128(3) release-prep analyse-trigger). (4) Land rock-solid fixes §11.4.123 + §11.4.43/.115 + §11.4.9. (5) Author NEW tests of ALL supported types §11.4.27 + §11.4.85, anti-bluff + paired §1.1 mutation. (6) Rebuild (full, not module-only) + reflash CLEAN target §11.4.108. (7) RESTART full validation from last tag to now §11.4.40 — RESTART not resume — both/all devices parallel §11.4.103/.119, recorded §11.4.128, real physical proofs §11.4.5/.69/.107, no bluff. BINDS the existing release anchors for the huge-blocker case (STOP→fix-all→process-recordings→new-tests-all-types→rebuild→reflash→full-restart + restart-not-resume), citing not duplicating. Composes §11.4.4/.40/.42/.9/.27/.85/.102/.108/.114/.115/.123/.128/.103/.119. Classification: universal (§11.4.17). Gate CM-COVENANT-114-129-PROPAGATION (literal 11.4.129) + rec. gate CM-HUGE-BLOCKER-FULL-RESTART + paired §1.1 mutation.

Canonical authority: Constitution.md §11.4.129. Release blocker. No --resume-after-blocker / --spot-validate-after-fix / --skip-recording-analysis / --skip-new-tests / --module-only-after-blocker / --single-device-restart escape.

§11.4.130 — Post-remediation validate-the-fix-FIRST-after-redeploy (User mandate, 2026-06-06)

Forensic anchor — direct user mandate (2026-06-06): when a blocker found during release validation is fixed and a new artifact (rebuild / new flashing image / redeploy) is produced + the target reflashed, FIRST re-test the SPECIFIC last-failing features + validate the just-incorporated fixes BEFORE the broader / full validation.

When a blocker / critical failure found during release validation is FIXED and a new artifact is produced + the target reflashed / redistributed / updated, in order, NO shortcut: (1) re-test the SPECIFIC last-failing features FIRST (targeted guard tests for exactly the defects this fix addressed) BEFORE any broader / full-suite validation; (2) validate the just-incorporated fixes with real captured evidence — the §11.4.115 RED test flips GREEN at RED_MODE=0 on the new artifact AND the §11.4.108 runtime-signature verifies on the CLEAN target the redeploy produced (metadata-only/config-only/absence-of-error/grep-without-runtime PASS forbidden §11.4/.1; proof §11.4.5/.69/.107/.123); (3) ONLY after the targeted fix is CONFIRMED working proceed to the §11.4.40 full retest from last tag to now. Rationale: a first fix attempt may not work / be incomplete / regress again under the

new artifact — confirming FIRST catches fix-did-not-take immediately instead of hours later at the END of a full cycle (then restarting §11.4.129); cheap-confirmation-first is §11.4.82 applied to the post-blocker reflash. This is §11.4.46 specialised for the post-blocker-refs flash case + the targeted-confirmation phase that GATES §11.4.129's step-7 full-restart. Honest boundary §11.4.6: “probably took” ≠ “took”; a still-FAILing targeted re-test re-enters the §11.4.114/.115 isolate→RED→fix loop, never proceeds to the full cycle on a still-broken fix. Composes §11.4.4/.40/.46/.108/.114/.115/.123/.129/.82. Classification: universal (§11.4.17). Gate CM-COVENANT-114-130-PROPAGATION (literal 11.4.130) + rec. gate CM-FIX-FIRST-AFTER-REDEPLOY + paired §1.1 mutation.

Canonical authority: Constitution.md §11.4.130. Release blocker. No --skip-targeted-retest / --full-cycle-first / --assume-fix-took / --validate-fix-at-end / --skip-red-green-flip-on-new-artifact escape.

§11.4.131 — Standing session-resumption file mandate (User mandate, 2026-06-07)

Forensic anchor — verbatim user mandate (2026-06-07): “Make this markdown a standard file which will be written EVERY TIME when we need fresh session out of the box! It MUST BE always up to date and in sync so whenever new session is created all we have to do is just point to it!”

Every project MUST maintain a SINGLE canonical, always-current **session-resumption file** at a fixed, project-declared standard path (declared once per §11.4.35, never moved without a §11.4.66 operator decision) — the OUT-OF-THE-BOX entry point for any fresh session (creating a new session = ONLY point the new agent at this one file). §11.4.131 promotes §11.4.127 (PREPARE-on-demand) into a STANDING, version-controlled ARTIFACT, ALWAYS present + ALWAYS in sync. (A) Exists at the declared path at all times (literal path per §11.4.35, never silently moved). (B) (Re)written whenever a fresh session is needed OR the live state materially changes (new HEAD / build-artifact id / phase / device state / in-flight job / blocking decision) — the §12.10 trigger set; a stale file = §11.4.131 violation (same class as a §12.10 stale-CONTINUATION). (C) Carries SHORT + FULL §11.4.127 variants; points to .remember/remember.md + docs/CONTINUATION.md read FIRST + git fetch; embeds exact live-state anchors (HEAD, artifact checksums, device serials, in-flight PIDs+logs, evidence paths); states PHASE + NEXT + terminal goal; restates binding constraints (anti-bluff §11.4, no-force-push §11.4.113, exact version/naming, hardware gotchas); MOMENT-VALID never generic (§11.4.6). (D) §11.4.65 export (.html/.pdf siblings) + §11.4.44 revision header. (E) Given ONLY this file's path a fresh session fully resumes with zero additional context. Composes §12.10/.127/.65/.44/.6/.66/.126. Classification: universal (§11.4.17). Gate CM-COVENANT-114-131-PROPAGATION (literal 11.4.131) + rec. gate CM-SESSION-RESUMPTION-FILE-PRESENT + paired §1.1 mutation.

Canonical authority: Constitution.md §11.4.131. Release blocker. No --skip-resumption-file / --ephemeral-prompt-only / --stale-resumption-OK / --generic-template-OK escape.

§11.4.132 — Risk-ordered validation priority mandate (User mandate, 2026-06-07)

Forensic anchor — verbatim user mandate (2026-06-07): “We MUST ALWAYS first test and validate features, functionalities and fixes/changes that have been worked most recently, the ones which were most problematic, which have the most chance to crash or break again, the ones which have been re-opened the most times! Then, after we validate and verify all this with real (physical) proofs and hard evidence, with no false results and bluffs of any kind, we continue with all other existing tests in the test suites! This IS MANDATORY.”

Tests / validations / verifications MUST run in RISK-DESCENDING order — the highest-risk set FIRST, and ONLY AFTER that set is fully GREEN with real (physical) captured evidence does the remainder of the suite run. Closed risk factors, highest-first: (a) most-recently-worked; (b) historically most-problematic (longest defect history, most prior fixes/failures); (c) highest crash/break/regress likelihood (greatest blast radius/complexity/dependency surface); (d) most-reopened per §11.4.55 reopens-count. Each highest-risk item MUST pass with real (physical) captured evidence per §11.4.5/§11.4.69/§11.4.107 — no metadata-only / config-only / absence-of-error / grep-without-runtime PASS (§11.4/§11.4.1), no false results, no bluff (§11.4.6). ONLY AFTER the entire highest-risk set is GREEN with captured proof does the rest of the suite run; arbitrary-order or lower-risk-before-highest-risk-GREEN = §11.4.132 violation. REFINES/STRENGTHENS §11.4.130 (generalises “validate-just-fixed-first” to the full risk-ordered set) + §11.4.46 (adds explicit risk-ordering) + §11.4.42 (applies implementation-priority discipline to VALIDATION ordering). Classification: universal (§11.4.17) — consuming project supplies its recency/problematic-history/reopen-count sources (e.g. §11.4.93 DB reopens_count+last_modified) per §11.4.35. Composes §11.4.4/.5/.6/.7/.40/.42/.46/.50/.55/.69/.107/.130. Gate CM-COVENANT-114-132-PROPAGATION (literal 11.4.132) + rec. gate CM-RISK-ORDERED-VALIDATION-PRIORITY + paired §1.1 mutation.

Canonical authority: Constitution.md §11.4.132. Release blocker. No --skip-risk-ordering / --any-order-OK / --suite-order-fixed escape.

§11.4.133 — Target-System + hardware safety mandate (User mandate, 2026-06-08)

Forensic anchor — verbatim user mandate (2026-06-08): “Make sure that all changes we do to the System are ALWAYS safe for the System itself and for the hardware the system runs on! This is MANDATORY.”

Every change to the TARGET system (firmware, kernel, init/boot scripts, drivers, sysfs/devfreq/voltage/clock/thermal/regulator register writes, partition/bootloader/U-Boot, HAL, framework, prebuilts, device config) MUST ALWAYS be safe for BOTH (a) the target System itself — MUST NOT brick, boot-loop, corrupt data, or render the device unrecoverable — AND (b) the hardware it runs on — MUST NOT exceed safe electrical/thermal/voltage/clock limits or damage panels/storage/radios/regulators. Concrete obligations: (1) reversible-first — verify irreversible high-blast-radius changes (bootloader/U-Boot MD5, partition layout) against known-good values + capture a pre-op backup (§9.2) BEFORE applying; (2) NO unverified hardware-control writes — never write an unverified value to a voltage/clock/regulator/thermal-throttle/current-limit sysfs node or register that could exceed

datasheet limits, the safe range established as FACT (§11.4.6), never guessed; (3) thermal/perf changes (forcing a performance governor, pinning the top OPP, disabling thermal management) MUST respect the device's cooling design, validated by captured thermal evidence; (4) flashing MUST use the sanctioned tool + a freshly-built integrity-verified image — never an ad-hoc partition write or stale/unverified artifact; (5) unprovable-safety ⇒ blocked — a change whose target/hardware safety cannot be established from captured evidence is treated as UNSAFE and blocked (§11.4.6 + §11.4.101 reversible-first + §11.4.123 rock-solid-proof). DISTINCT from §12 host-session safety: §12 protects the DEVELOPER's HOST + session; §11.4.133 protects the TARGET device + its hardware — both apply, neither weakens the other. Classification: universal (§11.4.17) — consuming project supplies its concrete hardware-control surfaces, datasheet-safe ranges, known-good bootloader/image hashes, and sanctioned flashing tool per §11.4.35. Composes §12 / §11.4.6 / §11.4.101 / §11.4.108 / §11.4.123. Gate CM-COVENANT-114-133-PROPAGATION (literal 11.4.133) + rec. gate CM-TARGET-HARDWARE-SAFETY + paired §1.1 mutation.

Canonical authority: Constitution.md §11.4.133. Release blocker. No --unsafe-hardware-write / --skip-system-safety / --brick-risk-accepted escape.

§11.4.134 — Code-review iterate-until-GO + rock-solid-evidence mandate (User mandate, 2026-06-08)

Forensic anchor — verbatim user mandate (2026-06-08): “For any fixes/changes given back to us for re-work by the code-review process, once we fix/improve everything per the code-review's requests, we MUST RE-RUN code-review AGAIN until we get a GO from it with NO new issues reported or warnings of any kind! All results produced by this whole process MUST ALWAYS give us rock-solid PHYSICAL evidence that the fixed/improved codebase really works now as expected, with no false results and no bluff(s) of any kind.”

When the §11.4.125 code-review returns ANY finding (BLOCKING, nit, or warning) and the author fixes/improves per that review, the code review MUST BE RE-RUN — and KEEP being re-run after each remediation round — until a clean GO with ZERO new issues AND ZERO warnings. A single “addressed the findings” pass is NOT sufficient: the corrected batch MUST pass a FRESH adversarial review (a re-review can surface NEW findings introduced by the fixes — the §11.4.1 fix-A-creates-B mode). The loop terminates ONLY on a clean GO; a residual warning is itself a finding that re-arms the loop. Every round's verdict AND every fix's validation MUST carry rock-solid PHYSICAL captured evidence per §11.4.5 / §11.4.69 / §11.4.107 (captured audio / video / sysfs / dumphsys / sink-side / runtime-signature) proving the fixed codebase REALLY works — never metadata-only / config-only / absence-of-error / grep-without-runtime; no false results, no bluff; a GO unbacked by physical evidence is a §11.4 PASS-bluff at the review-loop layer. REFINES / STRENGTHENS §11.4.125 (iterate-until-no-blocking): makes the loop EXPLICIT, raises termination to ZERO findings AND ZERO warnings, BINDS physical evidence to every round. Classification: universal (§11.4.17). Composes §11.4.125 / §11.4.1 / §11.4.4 / §11.4.5 / §11.4.6 / §11.4.69 / §11.4.107 / §11.4.50 / §11.4.108 / §11.4.123. Gate CM-COVENANT-114-134-PROPAGATION (literal 11.4.134) + rec. gate CM-CODE-REVIEW-ITERATE-UNTIL-GO + paired §1.1 mutation.

Canonical authority: Constitution.md §11.4.134. Release blocker. No `--skip-rereview` / `--single-review-pass` / `--warnings-ok` / `--evidence-optional` escape.

§11.4.135 — Standing regression-guard suite + every-fixed-defect-gets-a-permanent-regression-test (User mandate, 2026-06-08). Every project MUST maintain a STANDING regression-guard suite that runs on EVERY build+deploy and BLOCKS the release tag on any failure. Every closed defect (stable ticket id, e.g. ATM-NNN) MUST, in the SAME commit as its fix (extending the §11.4.43 DOCUMENT step), register a permanent §11.4.115 RED-on-broken-artifact regression test into the suite — `RED_MODE=1` capturing the historical defect on a prefix artifact (the proof the guard is real), `RED_MODE=0` the standing GREEN guard asserting the defect is ABSENT. A closure without a registered guard is a §11.4.123 violation. The suite runs FIRST in the post-deploy cycle (highest-risk set per §11.4.132) and is a §11.4.40 release-gate blocker. Forensic anchor (FACT): the wrong-subtitle-on-2nd-display defect was “fixed” via a source-side `CONTROL_MENU_LABEL_DENYLIST` that NO test mirrored or re-ran, so the NEXT chrome class recurred silently while the GREEN suite passed. Industry-standard bug-driven testing (Google content-driven testing; AOSP CTS/Tradefed) made mechanical + enforced. Composes §11.4.4 / §11.4.40 / §11.4.43 / §11.4.46 / §11.4.50 / §11.4.107 / §11.4.108 / §11.4.115 / §11.4.118 / §11.4.123 / §11.4.124 / §11.4.130 / §11.4.132. Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-135-PROPAGATION` (literal 11.4.135) + recommended gates `CM-REGRESSION-GUARD-REGISTERED` / `CM-REGRESSION-GUARD-SUITE-WIRED` + paired §1.1 mutation. **Canonical authority:** constitution submodule Constitution.md §11.4.135. Non-compliance is a release blocker. No escape hatch — no `--skip-regression-guard`, `--no-guard-on-close`, `--guard-optional` flag.

§11.4.136 — Real-content end-to-end playback-test mandate (User mandate, 2026-06-08). Refines/strengthens §11.4.107. Any test asserting media playback works MUST drive REAL content (catalog stream or offline reference clip) through the user’s path (§11.4.48 UI-driven → §11.4.117 CV/OCR fallback) and assert it genuinely PLAYS via the §11.4.107 liveness battery PLUS a decoder-health census — a numeric drop-buffer budget, no buffer-timestamp re-order/discard, no codec-reject (cite Android/Media3 ExoPlayer OEM pre-OTA playback-test mandate: “too many dropped buffers” >25, “unexpected presentation timestamp”, “test timed out”). Metadata-only / launch-only / registration-only / single-frame / config-only PASS is forbidden (§11.4 / §11.4.1). A golden/reference clip corpus (BBC ExoPlayer testing samples) is the offline ground-truth. Composes §11.4.5 / §11.4.48 / §11.4.50 / §11.4.107 / §11.4.117 / §11.4.123 / §11.4.13 / §11.4.69. Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-136-PROPAGATION` (literal 11.4.136) + recommended gate `CM-REAL-CONTENT-PLAYBACK-TEST` + paired §1.1 mutation. **Canonical authority:** constitution submodule Constitution.md §11.4.136. Non-compliance is a release blocker. No escape hatch — no `--launch-proves-playback`, `--skip-decoder-health`, `--metadata-playback-pass` suffices flag.

§11.4.137 — Subtitle/caption content-correctness oracle + secure-display-proxy-honesty mandate (User mandate, 2026-06-08). Refines §11.4.117 + §11.4.107 + §11.4.112. Forensic anchor (FACT): tests tasked to “physically verify the 2nd-display subtitle” PASSED GREEN while subtitles did NOT show / showed WRONG — the streaming player surface is `FLAG_SECURE` so `screencap -d <secondary>` returns BLACK (autonomous PIXEL verification structurally impossible per §11.4.112), so the test fell back to the accessibility-scraped/

`persist.atmosphere.subdebug` proxy, and the proxy accepted a chrome/menu LABEL (Аудио и субтитры) as a valid subtitle because the prose floor accepted any multibyte prose and NO menu-label denylist + NO position/cadence check existed. The mandate: a subtitle-correctness test MUST classify the cue's *content class* — a present cue is NOT a correct cue. CHROME (FAIL) if a known control/menu label (closed multilingual deny-list MIRRORED from source, case-folded incl. non-ASCII), time/numeric chrome, not prose, OUTSIDE the lower safe-title band (CEA-708 9-anchor grid), OR STATIC across the window (real subtitle changes → ≥2 distinct prose cues, a metamorphic relation). DIALOGUE (PASS) only when prose + not-denied + not-chrome + position-ok + cadence ≥2 OR fuzzy-matches the SOURCE-extracted cue via normalized edit distance (§11.4.123 host ground truth). The oracle MUST be self-validated golden-good/golden-bad (§11.4.107(10)) and the deny-list MUST be verified present in the SHIPPED artifact (§11.4.108) — a source-green denylist with no test mirror + no artifact check is the exact recurrence pattern forbidden here. Secure-display honesty (§11.4.112): where FLAG_SECURE makes pixel verification impossible, the rock-solid autonomous proof is the player's caption telemetry + source-track presence + content-class oracle — NEVER a faked pixel “physical” pass; human-eye pixel confirmation is `operator_attended` (§11.4.52) with a tracked migration item. App-agnostic (keys off content class). Composes §11.4.3 / §11.4.5 / §11.4.6 / §11.4.107 / §11.4.108 / §11.4.112 / §11.4.115 / §11.4.117 / §11.4.123 / §11.4.13 / §11.4.69. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-137-PROPAGATION (literal 11.4.137) + recommended gate CM-SUBTITLE-CONTENT-CORRECTNESS-ORACLE + paired §1.1 mutation (strip the denylist/position/cadence check → golden-bad Аудио и субтитры PASSES → gate FAILs). **Canonical authority:** constitution submodule Constitution.md §11.4.137. Non-compliance is a release blocker. No escape hatch — no `--present-cue-is-correct`, `--skip-chrome-oracle`, `--length-heuristic-suffices`, `--pixel-pass-on-secure-display`, `--skip-position-check`, `--skip-cadence-check` flag.

§11.4.138 — Operator-escape => mandatory bluff-audit + permanent guard (User mandate, 2026-06-08). When the operator (or any out-of-band channel) finds a defect that the GREEN test suite passed, this is by definition a §11.4 PASS-bluff — it MUST trigger, before the fix is closed: (1) a §11.4.102 systematic-debugging pass to FACT-root-cause; (2) a bluff-audit identifying the EXACT assertion that should have caught it but didn't, cited to `file:line` (canonical example: `lib/subtitle_content_validation.sh:sub_is_prose()` returning TRUE for Аудио и субтитры); (3) a permanent §11.4.135 regression guard registered in the SAME commit as the fix, with its §11.4.115 RED capturing the operator-found defect; (4) the bluff-audit committed under `docs/research/<scope>/<defect>_bluff_audit/`. Closing an operator-found defect WITHOUT the bluff-audit + permanent guard is itself a §11.4 violation (the bluff that let it through is still live and the defect will recur). Composes §11.4 / §11.4.1 / §11.4.102 / §11.4.108 / §11.4.115 / §11.4.118 / §11.4.123 / §11.4.135. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-138-PROPAGATION (literal 11.4.138) + recommended gate CM-OPERATOR-ESCAPE-BLUFF-AUDIT + paired §1.1 mutation. **Canonical authority:** constitution submodule Constitution.md §11.4.138. Non-compliance is a release blocker. No escape hatch — no `--close-without-bluff-audit`, `--operator-find-is-just-a-bug`, `--skip-permanent-guard` flag.

§11.4.139 — Fresh-process clean-artifact runtime-signature mandate (User mandate, 2026-06-08). Refines §11.4.108. Before any post-deploy validation — ESPECIALLY a non-pixel proxy verification (the subdebug/accessibility-cue channel used for FLAG_SECURE displays) — the harness MUST assert `running-artifact == built-artifact`: the deploy yielded a CLEAN target (mutable-overlay/

userdata wiped) OR a pre-validation check proves no stale overlay shadows the deployed code (e.g. every guarded package — incl. the Presenter that emits the subtitle cue — resolves to the system partition, no per-user override). A stale shadow of the cue-emitting component (e.g. a Presenter APK predating the denylist) makes the proxy report on code that was never deployed — any PASS is a §11.4 PASS-bluff. Each fix declares ONE machine-checkable runtime signature verified on the clean target (the §11.4.108 registry IS the definition of done); for the subtitle class the signature is “the shipped Presenter APK contains the denylist literal (case-insensitive) AND the subdebug channel emits candidate REJECTED reason=chrome-label for a menu label.” Composes §11.4.46 / §11.4.108 / §11.4.130 / §11.4.135 / §11.4.137. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-139-PROPAGATION (literal 11.4.139) + recommended gate CM-CLEAN-ARTIFACT-RUNTIME-SIGNATURE + paired §1.1 mutation. **Canonical authority:** constitution submodule Constitution.md §11.4.139. Non-compliance is a release blocker. No escape hatch — no --validate-against-running-state, --skip-clean-precondition, --shadow-OK flag.

§11.4.140 — Universal action-prefix system (ACTION_NAME ::) (User mandate, 2026-06-09; GRAMMAR_ADDENDUM 2026-06-09). When a user prompt’s FIRST non-blank line starts with a recognised action prefix, you MUST: (1) look the action token up in the action registry constitution/actions/registry.yaml (or \$HELIX_ACTION_REGISTRY); (2) if it is a registered action, REPLACE the prefix with that action’s expansion text and apply its rules; (3) execute the remainder of the prompt under the expanded instruction. **Four EQUIVALENT forms** — same action, same expansion, same execution: (1) ACTION_NAME :: <rest> (bare ::), (2) PREFIX::ACTION_NAME :: <rest> (namespaced ::), (3) / ACTION_NAME <rest> (bare slash), (4) /PREFIX::ACTION_NAME <rest> (namespaced slash). Thus BACKGROUND :: x ≡ DEFAULT::BACKGROUND :: x ≡ /BACKGROUND x ≡ /DEFAULT::BACKGROUND x. PREFIX is an action NAMESPACE; the reserved default namespace is **DEFAULT**, and an action runs WITH or WITHOUT the prefix. Grammar (all hold): anchored at the FIRST non-blank line only (mid-prose tokens never match); the action token AND the namespace are UPPERCASE-only [A-Z] [A-Z0-9_]* (lowercase never matches); the namespace separator :: inside the token carries NO surrounding spaces (PREFIX::ACTION_NAME), DISTINCT from the action-body separator " :: " (one ASCII space on each side of :: — avoids C++ Foo::Bar, YAML key: value, URLs) in forms 1/2 and the slash-body separator (one space) in forms 3/4; stacked prefixes (A :: B :: rest) apply outer-to-inner, left-to-right (expand A, re-scan, expand B, then the residual is the task); a leading \ escapes the prefix for BOTH the :: and the slash form (\BACKGROUND :: x, \ /BACKGROUND x — treat literally, strip the backslash, NO expansion) so action names can be discussed. **Conflict rule (slash form):** /ACTION_NAME (form 3) is honored as the action ONLY when ACTION_NAME (case-folded) does not collide with a built-in/host slash command (registry slash_bare: auto + slash_conflicts: [..]); form 4 (/PREFIX::ACTION_NAME) is ALWAYS unambiguous and always honored. An unknown token that matches the grammar shape (any of the 4 forms) but is NOT registered is NEVER silently expanded or silently dropped — ask which registered action was meant (§11.4.66 / §11.4.105) or treat it as a literal prompt, NEVER invent an expansion (§11.4.6); any prompt not satisfying the grammar is an ordinary prompt and the system is a no-op. The registered action **BACKGROUND** expands to: “*The following prompt that we will provide MUST BE executed in background in parallel with all main work*”

streams using the subagents-driven development approach! All work done MUST PRODUCE rock solid evidence covered with hard physical proof(s) that all done is working as expected and as specified without any false results and without any bluff!" (composes §11.4.20 / §11.4.70 subagent-driven, §11.4.58 / §11.4.103 parallel streams, §11.4.89 background execution, §11.4.5 / §11.4.69 / §11.4.107 captured physical evidence, §11.4 anti-bluff). The system is UNIVERSAL (every CLI agent reads this block via its context carrier per §11.4.35), extensible (new action = new registry row), decoupled + reusable (§11.4.28), and loads out-of-the-box. Classification: universal (§11.4.17). **Canonical authority:** constitution submodule Constitution.md §11.4.140. Non-compliance is a release blocker. No escape hatch — no --skip-action-prefix, --ignore-prefix, --no-registry, --invent-expansion-OK, --single-layer-only flag.

§11.4.141 — Token-efficiency mandate (research-derived + operator mandate, 2026-06-09). Every project worked on by AI coding agents MUST cut token spend (input AND output) toward **30–40% of current (a 60–70% reduction)** WITHOUT degrading quality/performance/safety or breaking any existing mechanism, via a composable, safety-ranked measure set: (1) **prompt-cache the static governance prefix** — the always-loaded governance forms a byte-stable cache-breakpointed prefix with no volatile bytes ahead of it; cache reads cost $\sim 0.1 \times$ base input (the dominant cost driver — measured $\sim 170K$ tokens of governance re-sent every turn, externally corroborated by Claude Code issue #24147); caching is transparent so it removes no rule, weakens no gate, changes no verdict — only billing (PRIMARY, biggest + safest lever); (2) **subagent model-tiering + output-to-file** — mechanical non-judgment work (search/grep/status/doc-export/read-only probes) to a Haiku-class model, the strong model RESERVED for all reasoning/verdicts/fix-design (§11.4.102)/code-review (§11.4.125)/demotion (§11.4.7), large output persisted to a file not an inline 350–520K-token transcript; the cheap model never emits a PASS so §11.4.50 + anti-bluff are untouched; (3) **thin always-loaded INDEX + on-demand detail** — concise index (one line per fix/anchor, EACH carrying the literal 11.4.N token so propagation gates pass) with the canonical full text kept gate-scanned in constitution/Constitution.md and reachable in one hop — a de-duplication realising §11.4.35, never a deletion; (4) **CodeGraph/retrieval-first over full-file loading** (§11.4.78/§11.4.79); (5) **output-token reduction** — terse status + effort: "low" on the mechanical allowlist only; (6) **tool-call batching + no re-reads**; (7) **compaction/context-editing for long sessions. Mandatory measured proof:** a token-accounting harness measures tokens-per-development-cycle BEFORE vs AFTER on a frozen deterministic workload from the authoritative usage object (input/cache_read/cache_creation/output split; NEVER tiktoken, NEVER the client-side cost estimate), reproduced N times (§11.4.50), pass = AFTER $\leq$ 40% of BEFORE OR the measured best-safe reduction with a cited cold-cache reason; the AFTER run MUST show ZERO regression on the pre-build sweep + meta-test mutation sweep + propagation gates + a strong-model reasoning probe + a cache-warm proof (cache_read_input_tokens > 0) — cost reduction with quality regression is a §11.4 FAIL. The headline number is the *measured* reduction, never the design estimate (§11.4.6/§11.4.123). No measure may break/degrade any existing mechanism, and the rule is structured so none can. Composes §11.4.5/.6/.20/.40/.50/.58/.69/.70/.78/.79/.80/.103/.106/.123/.125/.128/§12.6/§1.1. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-141-PROPAGATION (literal 11.4.141) + recommended gate CM-TOKEN-EFFICIENCY + paired §1.1 mutation (inject a pre-breakpoint volatile token → cache collapses → measured reduction falls below bar → gate FAILs). **Canonical authority:** constitution submodule Constitution.md §11.4.141. **Project instantiation**

(§11.4.35): [ATMOSphere consumer fills in: confirm the Claude Code governance-prefix cache is warm + free of pre-breakpoint volatiles; tier the §11.4.96-SAFE mechanical subagents (search/grep/status/doc-export) to Haiku with output-to-file under qa-results/; replace the consumer CLAUDE.md 112-row Applied-Fixes inline table + verbatim anchor restatements with a literal-anchor index pointing at constitution/Constitution.md; keep CodeGraph/lumen-first navigation; run the harness BEFORE/AFTER to prove the measured reduction.] Non-compliance is a release blocker. No escape hatch — no --skip-token-efficiency, --no-cache-governance, --assert-reduction-without-measuring, --tier-down-reasoning, --inline-all-governance, --tiktoken-estimate-OK flag.

§11.4.142 — Universal code-review mandate — every change reviewed, always, no exception (User mandate, 2026-06-09). Verbatim operator mandate: “ALL changes we do MUST pass through the code review step!!! ALWAYS!!!” EVERY change made to ANY repository this Constitution governs — without exception — MUST pass through an INDEPENDENT code-review step BEFORE it is accepted, committed, or built. NO change class is exempt: source code, fixes, tests, gates, meta-test mutations, documentation, doc-tooling, build/CI scripts, configuration, governance files (Constitution / CLAUDE / AGENTS / QWEN), conductor main-stream edits, sub-agent output, refactors, single-line edits — if a diff exists, it gets an independent review. This is the ABSOLUTE form of §11.4.125 (code-review agent gate after a batch, before pre-build sweep + main build): §11.4.142 strips every implicit scoping §11.4.125’s “after all fixes/changes/implementations are done” phrasing could leave open — no “just a doc edit”, no “just a one-liner”, no “the author already self-reviewed (§11.4.92)”, no “trivial change” carve-out.

Independence is load-bearing — the reviewer MUST be structurally separated from the author (a dedicated code-review agent, subagent-driven by default per §11.4.70 / §11.4.20, or a distinct human), NEVER the author re-reading their own work; §11.4.92’s multi-pass self-evaluation is the AUTHOR-side discipline and PRECEDES (never satisfies) §11.4.142. **The review is itself anti-bluff** (§11.4 / §11.4.1) — a rubber-stamp “looks good” is not a review; it MUST genuinely analyse correctness, safety (no host §12 / data §9 / target-hardware §11.4.133 regression), will-it-really-work (no solve-A-create-B), end-user behaviour (§11.4 / §107), and test-genuineness (§1.1), its conclusions captured evidence per §11.4.5 / §11.4.69, and **it iterates to a clean GO per §11.4.134** — any finding (BLOCKING / nit / warning) re-arms the loop, acceptance only on ZERO new findings + ZERO warnings with rock-solid physical evidence. Honest boundary (§11.4.6): a passing review does NOT replace §11.4.108 four-layer runtime-signature verification nor the §11.4.40 full-suite retest — it is one of MULTIPLE STRONG LAYERS, and the FIRST one every change crosses. Composes §11.4.1 / §11.4.4 / §11.4.5 / §11.4.6 / §11.4.20 / §11.4.40 / §11.4.69 / §11.4.70 / §11.4.92 / §11.4.108 / §11.4.110 / §11.4.125 / §11.4.134 / §107 / §1.1. Classification: universal (§11.4.17) — the consuming project supplies its reviewer-dispatch mechanism + change-acceptance seam (commit wrapper / merge queue / PR gate) per §11.4.35. Propagation gate CM-COVENANT-114-142-PROPAGATION (literal 11.4.142) + recommended gate CM-EVERY-CHANGE-REVIEWED (every accepted change carries a fresh independent-review-completed marker for its diff, produced by a reviewer distinct from the author, before acceptance/commit/build) + paired §1.1 mutation (strip the literal → propagation gate FAILs; accept a diff with no independent-review marker → CM-EVERY-CHANGE-REVIEWED FAILs). **Canonical authority:** constitution submodule Constitution.md §11.4.142. Non-compliance is a release blocker. No escape hatch — no --skip-review, --no-review, --trivial-change-exempt, --doc-edit-exempt, --self-review-suffices, --review-after-commit flag.

§11.4.143 — Real-user-journey mandate for video-streaming-app full-automation tests (User mandate, 2026-06-10). Verbatim operator mandate: “All video streaming apps ... require to choose some title and to press proper UI button to start or resume playing! Proper UI interaction to play exact show with proper content and subtitles is MANDATORY! Without it we just eventually play on 2nd display sample, and that’s it mostly! THIS MUST BE ADDED as MANDATORY RULE regarding testing of any video streaming app in general with full automation tests! Universal in root constitution + ATMOSphere project extensions.” Any full-automation test that asserts a video player / streaming application plays content MUST drive the REAL end-user journey through the app’s OWN UI — launch → BROWSE the actual catalog → choose a SPECIFIC title → press the real Play/Resume button → confirm THAT chosen content is genuinely playing, with its correct subtitles, on the intended routing target. A test that bypasses the journey with a sample / demo / built-in-loop clip, a deep-link / am start -a VIEW / intent shortcut, a synthetic or pre-staged stream, or any path that does NOT exercise the app’s own browse-select-play UI is a §11.4 PASS-bluff at the user-journey layer: it validates ROUTING (that *something* reaches the display) while leaving the user-visible behaviour the operator cares about — “the show I picked actually plays” — unproven (the operator’s “we just eventually play on 2nd display sample, and that’s it mostly” gap). The mandate (ALL hold): (1) **Real journey, not a shortcut** — launch → catalog browse → specific-title selection → real Play/Resume press through the app’s own UI (§11.4.48 UI-driven), NEVER a deep-link / intent / sample / loop-clip shortcut; (2) **Chosen-content confirmation** — the PASS proves the SPECIFIC selected title is playing (not merely that pixels move on the target) via the §11.4.107 liveness battery on the §11.4.136 real-content path + the §11.4.137 subtitle content-correctness oracle for the chosen title’s captions; (3) **Non-introspectable UIs use the pixel oracle** — when the app’s accessibility hierarchy is blank/unreliable (TV-Compose / leanback / canvas / GL), DRIVE input + ASSERT content via the §11.4.117 CV/OCR pixel oracle, never a hierarchy-only tool; (4) **Login via the credential single-source** (§11.4.10), never hardcoded, never logged; (5) **Honest SKIP, never a faked PASS** — where the autonomous journey is genuinely infeasible (hard human-only login / CAPTCHA, geo-block per §11.4.3, secure-surface pixel-blanking per §11.4.112) the test is operator_attended SKIP-with-reason per §11.4.52 + §11.4.3 with a tracked migration item — NEVER a metadata-only / sample-played / routing-only PASS. Honest boundary (§11.4.6): “the routing fired so the title is playing” is a guess — only the chosen-content liveness + subtitle oracle on the real journey proves it. Composes §11.4.48 / §11.4.107 / §11.4.117 / §11.4.136 / §11.4.137 / §11.4.52 / §11.4.3 / §107 / §1.1. Classification: universal (§11.4.17) — the consuming project supplies its concrete app roster, login-credential source, routing target, and UI-driving / pixel-oracle harness per §11.4.35. Propagation gate CM-COVENANT-114-143-PROPAGATION (literal 11.4.143) + recommended gate CM-VIDEO-REAL-JOURNEY-TEST (every video-streaming-app playback test drives the real browse-select-play journey + confirms the chosen content + subtitles, or SKIPS-with-reason) + paired §1.1 mutation (replace a real-journey test’s browse-select-play path with a deep-link / sample shortcut → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule Constitution.md §11.4.143. Non-compliance is a release blocker. No escape hatch — no --sample-playback-ok, --skip-real-journey, --deep-link-suffices, --routing-only-pass, --no-title-selection flag.

§11.4.144 — Tracked/recorded-device availability-following mandate (User mandate, 2026-06-10). Direct operator mandate: every device the project is tracking / following / recording (test / debug / manual-testing device, across every

reachable transport — USB / wireless ADB / SSH / serial / network introspection API) MUST be availability-FOLLOWED — its connection state continuously monitored, any drop handled, never silently abandoned and never presented as a continuous recording. The §11.4.128 always-on recorder KNOWS a tracked device is absent (its per-device loop guards on the reachable state) but, lacking a following discipline, merely spins idle — no data captured, no offline event logged, no resume, no escalation — so the recording corpus presents a continuous timeline with a silent hole, a §11.4 PASS-bluff at the recording-integrity layer (it claims continuous capture while no data exists for the gap). On a tracked device leaving its reachable state the system MUST, automatically: (a) **DETECT** the drop and **log an honest offline event** into the recording corpus (§11.4.6 — a silent gap presented as continuous capture is a fabricated-continuity bluff; §11.4.128 — a silent recording gap defeats always-on recording); (b) **WAIT** for the device to return using the project’s ALREADY-DEFINED reconnection timings — never invented numbers (§11.4.6), the SAME grace / reconnect / poll budgets the project’s recovery path already uses; (c) **RE-ATTACH** — resume recording / tracking the moment the device returns and log an honest online / resume event; (d) **ESCALATE** to the project’s sanctioned device-recovery path (per §11.4.69 feature class `device_recovery`) if the device does not return within the defined timeout — through the sanctioned, authorization-gated recovery entry point ONLY, never bypassing its gate and never performing a destructive recovery (e.g. a power-cycle) autonomously without that authorization (§11.4.21 / §11.4.101 — high-blast-radius recovery is gated; while blocked the system keeps following the device, logging the blocked-escalation honestly). A tracked-device drop that produces a silent corpus hole, a never-resumed recording, or a never-escalated permanent absence is the bluff this anchor forbids. Composes §11.4.128 (always-on recording — §11.4.144 closes its drop-handling gap) / §11.4.69 (`device_recovery` sink-side positive evidence) / §11.4.6 (honest offline / online events, reused-not-invented timings) / §11.4.14 (watchdog children reaped on stop) / §11.4.21 + §11.4.101 (gated, non-autonomous escalation). Classification: universal (§11.4.17) — the consuming project supplies its concrete tracking transport, device set, already-defined reconnection timings, and sanctioned recovery entry point per §11.4.35. Propagation gate CM-COVENANT-114-144-PROPAGATION (literal 11.4.144) + recommended gate CM-DEVICE-AVAILABILITY-FOLLOWED + paired §1.1 meta-test mutation (strip the watchdog wiring / let an absence go unlogged → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule Constitution.md §11.4.144. Non-compliance is a release blocker. No escape hatch — no `--skip-availability-following`, `--silent-recording-gap-OK`, `--no-reconnect-wait`, `--invent-reconnect-timing`, `--skip-recovery-escalation`, `--autonomous-power-cycle-OK` flag.

§11.4.145 — Independent multi-angle impact-research per change (User mandate, 2026-06-10). For EVERY fix / change / new feature, INDEPENDENT impact-research agents (subagent-driven §11.4.70/§11.4.20, structurally separate from the author — §11.4.92 self-eval PRECEDES, never satisfies — adversarial “refute-the-change” stance to defeat the documented LLM confirmation-bias failure mode) MUST research the change AND its connected/dependent features across a CLOSED SET OF EIGHT ANGLES — (1) correctness/logic, (2) regression (what existing working feature could break — via call-graph impact enumerating every direct + transitive caller and contract-dependent feature, each mapped to its test), (3) latent/dangerous-code — the recents class (runtime-only failure, interface/ABI/contract mismatch, concurrency/data-race, resource lifecycle), (4) security (taint/injection/secret/unsafe-API), (5) performance (latency/memory/thermal under load vs baseline), (6) host/data/target-hardware safety per §12/§9/§11.4.133, (7) cross-

feature interaction (shared state/timing/hardware contention), (8) business-logic conformance (matches the spec AND the connected features' contracts, not merely compiles). Forensic anchor (FACT, project-internal): the recents / 3-button-nav path shipped a latent AIDL-interface-contract mismatch that PASSED code review yet failed at RUNTIME ONLY (ANR on a recents-reaping gesture) because no review angle interrogated the interface contract / concurrency-lifecycle / silently-broken connected feature — the §11.4.108 SOURCE→ARTIFACT→RUNTIME→USER-VISIBLE gap caught only after the fact; §11.4.145 shifts that discovery LEFT. Each angle MUST name the TOOL(S) used + obtain the required DEPTH (the diff, the full changed unit, its declared contract, the spec section, every connected feature — NEVER the diff lines alone) + emit a captured-evidence conclusion per §11.4.5/ §11.4.69 (“looks fine” forbidden, §11.4.1); a genuinely-N/A angle is recorded NOT_APPLICABLE: <reason> per §11.4.6, never silently skipped. Output = a per-change impact-research REPORT (one section per angle: tool + evidence path + verdict) + a single GO/NO-GO that BLOCKS acceptance/commit/build on ANY unmitigated risk in ANY angle; the change is fixed/mitigated + affected angles re-researched, iterating to a CLEAN GO (zero unmitigated risk + zero warning) per §11.4.134 before the §11.4.125 review gate. Honest boundary (§11.4.6): a GO proves cross-angle internal-consistency + bounded blast-radius — it does NOT replace §11.4.108 runtime-signature verification nor §11.4.40 full-suite retest; it is one of MULTIPLE STRONG LAYERS, the FIRST research pass every change crosses. Composes/strengthens §11.4.92/.125/.108/.110/.134/.78/.79/.107/.85/.133/ §12/§9/.99/.6/§1.1. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-145-PROPAGATION (literal 11.4.145) + recommended gate CM-IMPACT-RESEARCH-PER-CHANGE + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; accept a change with a report missing an angle or an unmitigated NO-GO → CM-IMPACT-RESEARCH-PER-CHANGE FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule Constitution.md §11.4.145. Non-compliance is a release blocker. No escape hatch — no --skip-impact-research, --single-angle-suffices, --author-researches-own-change, --diff-skim-OK, --no-go-overridable, --trivial-change-no-research flag.

§11.4.146 — Reproduce-first test + same-test-confirms-fix + mandatory extend-to-all-cases workflow (User mandate, 2026-06-10). Every reported problem MUST be handled by a NAMED three-step test workflow — §11.4.146 does NOT re-author its component disciplines, it NAMES + BINDS them into the operator's exact per-defect sequence and adds ONLY the two emphases the components leave implicit. **(STEP 1 — REPRODUCE-FIRST + INVESTIGATE)** before any fix, author the §11.4.43 RED test as a §11.4.115 RED-baseline-on-the-broken-artifact (reproduce the defect on the CURRENT pre-fix artifact, capture defect-present physical evidence per §11.4.5/§11.4.69/§11.4.107); NEW EMPHASIS (D1) that same RED test is ALSO a deliberate investigation instrument per §11.4.102(A) — it MUST gather ADDITIONAL forensic data characterising the defect (triggers, boundaries, input/topology scope, adjacent failure modes) that feeds BOTH the fix design AND the STEP-3 extend scope; a RED test used only as a binary present/absent gate with no characterisation satisfies §11.4.115 but NOT §11.4.146 STEP 1. **(STEP 2 — SAME-TEST-CONFIRMS-FIX)** after the fix the SAME test source (the §11.4.115 polarity switch flipped RED_MODE=1→0) confirms the defect ABSENT — RED-on-broken then GREEN-on-fixed, both captured, on a clean target per §11.4.108/§11.4.139, validated-first-after-redeploy per §11.4.130, deterministic per §11.4.50; no separate happy-path test substitutes (the §11.4.43/§11.4.115 PASS-bluff). **(STEP 3 — EXTEND-TO-ALL-CASES, mandatory per-fix)** NEW EMPHASIS (D2) immediately after STEP 2 the test MUST be FANNED OUT

across the full case-space of the SAME functionality — all flows, valid + invalid + boundary (§11.4.85 empty/max/off-by-one) + concurrent (§11.4.85 contention) + failure-injection (§11.4.85 chaos) + topology variants (§11.4.3) — confirming no issue of any kind, with PROVABLE enumerated coverage per §11.4.118 (a listed case-set with per-case outcome, never “no other issues found”); a REQUIRED step, NOT deferred to a release-cycle discovery sweep and NOT reduced to a single guard; each user-visible case carries rock-solid physical evidence per §11.4.123 + is registered into the §11.4.135 standing regression-guard suite; a newly-discovered fan-out defect triggers §11.4.4 test-interrupt + re-enters STEP 1. (ANTI-BLUFF, all steps) every PASS is rock-solid CAPTURED physical evidence per §11.4.123 (§11.4.5/§11.4.69/§11.4.107) — metadata-only / config-only / absence-of-error / grep-without-runtime PASS forbidden (§11.4/§11.4.1); unclear validation method ⇒ deep-research-before-declaring-untestable per §11.4.123/§11.4.8/§11.4.99; an operator-found defect the green suite missed triggers §11.4.138. Honest boundary (§11.4.6): STEP 3 reduces the unknown-unknown surface but does not prove zero remaining defects (§11.4.118) — un-exercised cases stated as honest gaps, never silently implied clean. Composes §11.4.43/.115/.102/.130/.108/.139/.50/.85/.118/.135/.3/.123/.5/.69/.107/.138/.4/§107/ §1.1. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-146-PROPAGATION (literal 11.4.146) + recommended gate CM-REPRODUCE-FIRST-THEN-EXTEND (every closed defect’s fix carries a §11.4.115 reproduce-first polarity test with captured defect-characterisation [STEP 1] + its RED_MODE=0 GREEN confirmation [STEP 2] + an enumerated per-functionality extend case-set registered into the §11.4.135 suite [STEP 3]; a closure with the reproduce→confirm pair but NO enumerated extend case-set FAILs) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; close a defect with only the reproduce→confirm pair and no extend-case-set → CM-REPRODUCE-FIRST-THEN-EXTEND FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule Constitution.md §11.4.146. Non-compliance is a release blocker. No escape hatch — no --skip-reproduce-first, --fix-without-red, --skip-extend-to-all-cases, --reproduce-confirm-suffices, --defer-extend-to-release-sweep, --single-guard-suffices, --no-defect-characterisation flag.

§11.4.147 — Crashed-agent respawn-until-complete + no-work-loss registry mandate (User mandate, 2026-06-10). Verbatim operator intent: “any agent that crashed because of something MUST BE respawned and finish its work at some point! We MUST NOT lose any work, forget about or have it corrupted!” Forensic case study (FACT): dispatched subagents died on TRANSIENT causes (API Error: Server is temporarily limiting requests rate-limit killed 5 at once, API Error: socket connection closed unexpectedly killed 2, one left PARTIAL edits — two source files written, the dependent SystemUI sliders + doc + test NOT done), each re-dispatched by pure conductor vigilance with NO mechanical guarantee — the moment conductor context is lost / compacted / the conductor itself crashes, an in-flight-but-dead work unit is silently dropped (lost / forgotten / corrupted). Every dispatched agent/subagent MUST be tracked through its full lifecycle so a crash NEVER loses, forgets, or corrupts its work; a crashed agent is NOT a completed agent (§11.4.6 — “it probably finished before dying” is a forbidden guess; abnormal termination is positive evidence the unit is OPEN). The mandate (ALL hold): **(a) durable agent REGISTRY** — every dispatched agent has an append-only machine-readable registry entry (stable id, task + §11.4.58 file-scope, declared output path(s) + §11.4.5/§11.4.69 evidence dir, dispatch timestamp, status from the CLOSED SET {dispatched | in-flight | crashed | respawned | complete}), reusing the §11.4.116 sync substrate (JSONL event

stream + atomic status snapshot; “an entry with no dispatch-event cannot show complete”); the registry is the SINGLE SOURCE OF TRUTH for “is this work owed?” and an unregistered dispatch is itself a violation; **(b) mechanical CRASH-DETECTION + RESPAWN-until-complete** — any abnormal termination (transient rate-limit/socket-close, any non-completion exit, or a terminal error after the runtime’s own retries are exhausted) flips the entry to crashed + keeps the unit OPEN, and the unit is respawned (fresh agent claims the same id + scope + preserved state) and re-respawned until an agent reaches complete; respawn is the safe/reversible/bounded decision taken autonomously per §11.4.101 (never blocks the loop for a human), transient causes use the project’s ALREADY-DEFINED backoff budgets (never invented, §11.4.6), a deterministically-reproducible non-transient terminal error is investigated per §11.4.102 before further respawn (never a blind retry loop); **(c) PARTIAL-STATE preserve → §11.4.84-check → resume-or-clean-restart** — the crashed agent’s uncommitted edits + output doc are PRESERVED (never silently discarded → lost, never blindly committed → corruption), the respawn runs the §11.4.84 quiescence check on the preserved tree (every modified file accounted-for vs scope, no mutation// always pass/_mutated_* residue, no half-written/torn artifact), then EITHER RESUMES idempotently when it passes OR CLEAN-RESTARTS from a known-good base (§9.2 pre-op backup, reversible §11.4.101) when inconsistent — nothing lost, nothing corrupted; per-agent git worktree isolation (§11.4.58 L4 / §11.4.84) keeps the partial tree from contaminating other streams; **(d) “crash ≠ done” COMPLETION criterion** — a unit is DONE only when an agent reaches complete with its required captured evidence/output landed (§11.4.5/§11.4.69 + the §11.4.116 verdict-carries-evidence-path rule); the endless-loop done-condition (§11.4.87/§11.4.94/§11.4.97/§11.4.126) MUST NOT read satisfied while any entry is dispatched/in-flight/crashed/respawned-not-yet-complete, the zero-idle survey (§11.4.94) treats every non-complete entry as an OPEN item to reclaim, and a registry showing complete without landed evidence is a §11.4 PASS-bluff at the agent-lifecycle layer. Honest boundary (§11.4.6): the registry + respawn guarantee work is not lost/corrupted, NOT that it is correct — the respawned output still crosses §11.4.108 + §11.4.125/§11.4.142 review + §11.4.40 retest; §11.4.147 is the durability layer beneath those, the agent-side analogue of §11.4.144 (same detect→wait/backoff→re-attach/respawn→escalate shape). Composes §11.4.6/.58/.84/.87/.94/.97/.101/.116/.102/.108/.125/.142/.40/.128/.144/§9.2/§1.1. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-147-PROPAGATION (literal 11.4.147) + recommended gate CM-CRASHED-AGENT-RESPAWN-TRACKED + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; mark a crashed entry as the loop’s done-condition complete without landed evidence, OR drop a dispatched agent without a registry entry → CM-CRASHED-AGENT-RESPAWN-TRACKED FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule Constitution.md §11.4.147. Non-compliance is a release blocker. No escape hatch — no --skip-agent-registry, --crash-equals-done, --no-respawn, --discard-partial-state, --blind-commit-partial, --forget-dead-agent, --loop-done-with-crashed-entries flag.

§11.4.148 — Workable-item integrity (status+type+id) + comprehensive structured description + bidirectional external-tracker sync + BLOCKED unblock-choices mandate (User mandate, 2026-06-10). BINDS + STRENGTHENS the workable-item discipline (§11.4.15 status / §11.4.16 type / §11.4.54 id / §11.4.21 Operator-blocked / §11.4.91 description-clarity / §11.4.93 + §11.4.95 SQLite SSOT / §11.4.106 docs_chain) into ONE integrity contract spanning DB [?] docs [?] external tracker, adding the operator’s three emphases: **(D1) no item**

without a valid status + valid type + stable id — on ALL three surfaces (an item missing any of the three in the DB, the rendered docs, OR the external tracker FAILs the validator — release-blocker; the id is the cross-surface binding key); **(D2) comprehensive structured description per item** — WHAT it is (§11.4.91 ≥6-word/≥40-char clear meaning) + HOW it manifests + HOW to reproduce (§11.4.115/.146) + ACCEPTANCE CRITERIA (§11.4.123/.69 captured-evidence verdict), in the §11.4.93 DB description column AND the docs AND the tracker, never a stub/§11.4.91-anti-pattern fragment; **(D3) BLOCKED items carry WHY + enumerated unblock CHOICES** — tightens §11.4.21: every Operator-blocked item (incl. Blocked/BLOCKED documented alias normalised to canonical Operator-blocked, never a silent fork §11.4.6) MUST enumerate the closed list of decisions/actions that would unblock it ([A]...·[B]...·[C]..., mirroring §11.4.66's 2–4-option shape); a BLOCKED item with no enumerated choices FAILs; **(D4) regular never-missed bidirectional DB↔docs↔tracker sync** — the §11.4.93/.95 git-tracked SQLite DB is the SSoT, docs + tracker DERIVED; full sync (db-to-md / md-to-db byte-identical round-trip §11.4.93 + tracker push) runs regularly + on every change, §11.4.86 drift-proof fingerprint (sha256 of the sorted item keyset, NOT mtime) gates freshness, §11.4.106 docs_chain-bound so drift is mechanically caught not vigilance-dependent; **(D5) generic external-tracker push** carries statuses (collapsed onto the tracker's native set per §11.4.33/.112 when fixed, precise value preserved in a header, never lost), types, assignee (§11.4.104 participant handle, UNSET defaults from a project env var, never hardcoded/logged §11.4.10), and sub-tasks, IDEMPOTENT (dry-run-then-real, match-by-stable-key [<ID>] prefix / id custom-field, present⇒UPDATE/absent⇒CREATE, rate-limited, credential-redacted, sink-side created=N updated=M failed=0 proof §11.4.69), the MACHINERY project-agnostic §11.4.28 (consumer registers its tracker/list/field-map at runtime). Anti-bluff §11.4: every sync + validator pass carries captured evidence §11.4.5/.69; honest boundary §11.4.6 — guarantees well-formed items + agreeing surfaces, NOT that the underlying work is correct (still crosses §11.4.108/.40/.123). Composes §11.4.15/.16/.21/.33/.34/.54/.66/.86/.91/.93/.95/.104/.106/.112/.123/.10/.28/.69/.6/ §1.1. Classification: universal (§11.4.17) — consumer supplies its DB path, id prefix, tracker service + list/board id + field map + default-assignee env var, docs_chain context per §11.4.35. Propagation gate CM-COVENANT-114-148-PROPAGATION (literal 11.4.148) + recommended gates CM-ITEM-INTEGRITY-STATUS-TYPE-ID / CM-ITEM-COMPREHENSIVE-DESCRIPTION / CM-BLOCKED-UNBLOCK-CHOICES / CM-TRACKER-SYNC-IDEMPOTENT + paired §1.1 mutation (gate-code = separate work item). **Canonical authority:** constitution submodule Constitution.md §11.4.148. Non-compliance is a release blocker. No escape hatch — no --item-without-status, --item-without-type, --item-without-id, --stub-description-OK, --blocked-without-choices, --skip-tracker-sync, --one-way-sync-OK, --non-idempotent-tracker-push, --hardcode-assignee flag.

§11.4.149 — Per-workable-item testing-diary mandate (User mandate, 2026-06-10). Every workable item MUST carry an append-only TESTING DIARY — a chronological record of every test run against it — distinct from §11.4.93 item_history (lifecycle STATE transitions) + §11.4.55 Reopens (reopen cycles): the diary records TEST EXECUTIONS (which may or may not change status). The mandate (ALL hold): **(a)** a test_diary table additive to the §11.4.93/.95 SQLite SSoT, one row per run soft-keyed to the item id, capturing ISO-8601-UTC date_time + tested_by (closed set User|Operator|AI-agent|HelixQA) + result (§11.4.45 PASS|FAIL|SKIP + detail) + in-depth observations (long markdown, facts or explicit UNCONFIRMED:/PENDING_FORENSICS: §11.4.6, the

background a future fix needs) + action_taken+status_changed+from/to (did this run change status + WHY) + §11.4.69 evidence_path+feature_class, with a SCHEMA CONSTRAINT making a PASS row WITHOUT a non-empty evidence path impossible (a PASS-bluff rejected by the schema itself); **(b)** BOTH an in-depth diary doc + a derived at-a-glance summary VIEW (total/pass/fail/skip + last-verdict + last-run + status-changes + distinct testers/feature-classes, DERIVED never duplicated §11.4.93) feeding §11.4.132 risk-ordering; **(c)** four-format per-item Diary/Diary_Summary exports §11.4.65 + §11.4.86-drift-proof-fingerprinted (sha256 of the sorted diary keyset) + §11.4.106 docs_chain-bound so a diary row not exported/pushed FAILS a gate (never-missed); **(d)** external-tracker SUB-TASK model — the item task’s description stays the item (§11.4.148 D2, NOT polluted with diary text), each run a CHILD sub-task (type task) with a {TODO|In-progress|Completed} lifecycle (collapsed per §11.4.33/.112), observations + action + an Evidence: line (path not raw artefact §11.4.10/.13) + a diary-entry idempotency key, reusing the §11.4.148 D5 idempotent rate-limited credential-redacted sink-side-proven push, mapper project-agnostic §11.4.28; **(e)** MINIMAL-LLM deterministic bash/Go tooling (zero LLM in the data path — the observations prose is authored by whoever ran the test; tooling only stores/renders/pushes/validates), 100% test-covered §11.4.27 (unit + integration-against-the-real-tracker with an honest §11.4.3 SKIP-when-token-absent never a faked PASS + export + HelixQA Challenge + paired §1.1 mutation) so a diary PASS-bluff is mechanically impossible. Honest boundary §11.4.6 — the diary guarantees a complete auditable per-item test-history, NOT that any single run’s verdict is correct (rests on its own §11.4.69/.107/.123 evidence). Composes §11.4.6/.27/.45/.50/.55/.65/.69/.86/.93/.95/.106/.107/.123/.132/.148/.10/.13/.28/.33/.112/§1.1. Classification: universal (§11.4.17) — consumer supplies its DB path, diary doc layout, export formats, tracker sub-task field map, docs_chain context per §11.4.35. Propagation gate CM-COVENANT-114-149-PROPAGATION (literal 11.4.149) + recommended gates CM-TEST-DIARY-SYNC / CM-DIARY-PASS-REQUIRES-EVIDENCE + paired §1.1 mutation (gate-code = separate work item). **Canonical authority:** constitution submodule Constitution.md §11.4.149. Non-compliance is a release blocker. No escape hatch — no --skip-testing-diary, --diary-without-evidence, --no-diary-summary, --diary-without-tracker-subtask, --llm-in-diary-data-path, --diary-export-optional flag.

§11.4.150 — Mandatory deep multi-angle web research per change/issue, before declaring fixed or structural (User mandate, 2026-06-11). Verbatim operator mandate: “For every single issue we fix or improvement we make — besides tight systematic-debugging, fixing, code review by independent agents, comprehensive tests — we MUST ALWAYS do deep web research from various angles! No matter how big/small/simple/complex, dig deep the internet for articles, technical documentation, APIs, open-source code — EVERYTHING that can help make the best possible solution OR confirm we don’t have a more serious problem we’re unaware of! We MUST do everything possible to STOP the constant issue-reopening and finally start closing items as fixed+working! Do this ALWAYS in parallel with the main work stream!” For EVERY fix / improvement / change / closure — no matter how big, small, simple, or complex — the agent MUST, IN ADDITION to tight systematic-debugging (§11.4.102), the fix, independent-agent code review (§11.4.125 / §11.4.142), and comprehensive multi-layer tests (§11.4.4(b) / §11.4.40), perform a DOCUMENTED **deep multi-angle web research pass** digging the internet from VARIOUS angles — articles, official + vendor technical documentation, API references, standards, issue trackers, maintainer guidance, reusable open-source code — to BOTH (i) discover the best possible solution AND

(ii) confirm there is NOT a more serious underlying problem the team is unaware of. Forensic case study (FACT, 2026-06-11): a subtitle-on-secondary-display goal verdicted Won't-fix: structurally-impossible (§11.4.112 secure-surface pixel-blanking) was OVERTURNED by a deep multi-angle research pass that surfaced the real OCR-of-the-PRIMARY-display path — a “we can’t” became a shipped capability; the canonical anti-pattern of a structural verdict reached for want of research. The mandate (ALL hold): **(A) No closure-as-fixed/structural WITHOUT a documented deep-research pass** — no item may be marked Fixed/Implemented/Completed (§11.4.33) OR classified structurally-impossible won't-fix (§11.4.112) until a deep multi-angle pass is documented; §11.4.150 makes §11.4.8's pass UNCONDITIONAL (every item, however trivial, at the closure AND structural-verdict gates). **(B) Multiple angles, not a single lookup** — ≥ 2 genuinely-distinct angles (official-docs / standards / known-bug-trackers / alternative-approach / failure-mode / security / performance / platform-constraint / open-source-precedent) so a single-source confirmation-bias miss (§11.4.145) cannot pass; a one-link drive-by is NOT deep multi-angle. **(C) Confirm-no-bigger-problem, not just find-a-fix** — explicitly seek evidence the fix does not mask a deeper defect AND the closure/verdict is genuinely safe; “we found nothing worse” requires the enumerated search (§11.4.118). **(D) Latest-source + cited** — LATEST authoritative versions per §11.4.99 (never training-data/memory), each cited by URL + access date in the research artefact AND the closure commit footer (Deep-research <date>: <urls> OR the literal NO external solution found – original work per §11.4.8). **(E) Reopen-breaking is the PURPOSE** — STOP the constant Fixed[?]Reopened churn (§11.4.34 / §11.4.55); a reopen whose root cause a deep pass would have surfaced is a §11.4.150 miss; composes §11.4.7 (demotion-evidence) + §11.4.112 (structural verdict now REQUIRES the cited-authorities pass). **(F) ALWAYS in parallel with the main stream** — background subagent-driven (§11.4.70 / §11.4.20 / §11.4.103 / §11.4.89) concurrent with main fix/build/test work, NEVER serialising it or stalling the loop (§11.4.94 / §11.4.97 / §11.4.101 / §11.4.126). **(G) Apply ASAP to every workable item** — retroactively + going forward across the §11.4.93 / §11.4.95 SSoT; items closed without a documented pass are re-audited in the §11.4.40 / §11.4.42 release-gate sweep. Honest boundary (§11.4.6): the pass reduces the unknown-unknown surface + breaks the most common reopen causes — it does NOT prove zero remaining defects (§11.4.118) and does NOT replace §11.4.108 runtime-signature verification, §11.4.125 / §11.4.142 review, or §11.4.40 retest; it is the research layer every fix/closure/structural-verdict additionally crosses, and “we probably don’t have a bigger problem” without the enumerated multi-angle search is a guess (§11.4.6), never a finding. Composes §11.4.8 / §11.4.99 / §11.4.123 / §11.4.118 / §11.4.145 / §11.4.125 / §11.4.142 / §11.4.7 / §11.4.112 / §11.4.34 / §11.4.55 / §11.4.70 / §11.4.20 / §11.4.89 / §11.4.103 / §11.4.40 / §11.4.42 / §11.4.93 / §11.4.95 / §11.4.6 / §1.1. Classification: universal (§11.4.17) — the consuming project supplies its concrete research corpora, angle set, item tracker, and closure-commit-footer convention per §11.4.35. Propagation gate CM-COVENANT-114-150-PROPAGATION (literal 11.4.150) + recommended gate CM-DEEP-RESEARCH-PER-ISSUE (every closed/structural-verdicted item carries a documented multi-angle deep-research artefact + cited-source closure footer, run in parallel, before the closure/structural verdict is accepted) + paired §1.1 mutation (strip the literal → propagation gate FAILs; close an item or reach a structural verdict with no documented multi-angle research artefact / cited footer → CM-DEEP-RESEARCH-PER-ISSUE FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule Constitution.md §11.4.150. Non-compliance is a release blocker. No escape hatch

— no --skip-deep-research, --single-source-suffices, --trivial-change-no-research, --close-without-research, --structural-without-research, --serialise-research, --research-from-memory-OK flag.

§11.4.151 — Project-prefixed release-tag/version-naming mandate (User mandate, 2026-06-12). Verbatim operator mandate: “Every release tag and version name we create — on the main repository and on every Submodule we own — MUST be prefixed with the project’s release prefix, e.g. `helix_translate-1.0.0-dev-0.0.1`. The prefix MUST come from `HELIX_RELEASE_PREFIX` in our `.env` if it is set, otherwise from the lowercased project root directory name. The SAME prefix MUST be used across the main repo and all owned Submodules in one release so a release is greppable across every repository.” Every release tag AND every version name created on the main repository AND on every owned-by-us submodule (§11.4.28) MUST be prefixed with the project’s release prefix, form `<PREFIX>-<version>` (canonical example: `helix_translate-1.0.0-dev-0.0.1`), so a release is identifiable + greppable across every repository it spans (one `git tag -l '<PREFIX>-*'` enumerates the whole release surface). **Prefix resolution order (closed-set, deterministic — §11.4.6 no-guessing):** (1) `HELIX_RELEASE_PREFIX` from the project’s `.env` — authoritative when set; `.env` is git-ignored per §11.4.30 and the variable is documented in the tracked `.env.example` (a §11.4.77 re-obtain mechanism, never committed); (2) fallback = the lowercased snake_case form of the project root directory name (no spaces) per §11.4.29 — used whenever the env var is unset/empty, so a prefix is ALWAYS resolvable from the checkout with zero operator input. **The prefix MUST be IDENTICAL across the main repo and all owned submodules within a single release** — a release that tags the main repo `<PREFIX>-1.2.0` while tagging an owned submodule with a different/unprefixed value is a §11.4.151 violation (the cross-repo grep no longer enumerates the release). Version codes increment monotonically within the prefix (`<PREFIX>-...-0.0.1 → <PREFIX>-...-0.0.2 → ...`), never reset, never skipped. Honest boundary (§11.4.6): the prefix guarantees a release is identifiable + uniform across every repository, NOT that its contents are correct — the tag is still created only after the §11.4.40 full-suite retest GREEN and reaches every upstream via the §11.4.113 merge-onto-latest-main path (NEVER a force-push), fanned out per §2.1. Composes §2.1 / §11.4.29 / §11.4.30 / §11.4.40 / §11.4.113 / §11.4.126 (the release-scope terminal condition is a published, prefixed tag). Classification: universal (§11.4.17) — the consuming project supplies its concrete prefix value + the `HELIX_RELEASE_PREFIX` env var per §11.4.35. Propagation gate `CM-COVENANT-114-151-PROPAGATION` (literal 11.4.151) + recommended gate `CM-RELEASE-PREFIX-NAMING` (every release tag/version on the main repo + every owned submodule carries the resolved `<PREFIX>-` prefix, identical across the release) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; create an unprefixed or differing-prefix release tag → `CM-RELEASE-PREFIX-NAMING` FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.151. Non-compliance is a release blocker. No escape hatch — no --no-release-prefix, --unprefixed-tag, --prefix-optional, --differing-submodule-prefix flag.

§11.4.152 — Crashlytics-recorded-data continuous monitoring + systematic-debug + regression-test-coverage mandate (User mandate, 2026-06-13).

Verbatim operator intent: “For every project that has Firebase Crashlytics enabled / wired, we MUST continuously monitor ALL of the Crashlytics-recorded data — crashes, ANRs, performance traces, and non-fatals — systematically debug each, fix and improve, and cover everything with validation and verification tests. This MUST be checked regularly, with no false results and no bluff of any kind!” Every project that has Firebase Crashlytics enabled/wired (SDK linked into a shipping

artifact, crash+non-fatal+ANR reporting active) MUST treat the Crashlytics console as a first-class captured-evidence channel from real end-user devices and continuously process every datum it records — an open Crashlytics issue with a green test suite is a §11.4 PASS-bluff at the field-telemetry layer (a real user hitting a broken feature while the suite reports green). STRENGTHENS+COMPLETES §11.4.47: §11.4.47 owns the periodic REVIEW + dedup + Issue-creation pass that SURFACES Crashlytics/Analytics/Performance findings; §11.4.152 owns what happens to each surfaced item AFTER — the systematic-debug → fix/improve → regression-test-coverage lifecycle that drives it to a proven regression-immune closure (§11.4.47 finds it; §11.4.152 fixes it + proves it stays fixed; both mandatory, neither substitutes). The mandate (ALL hold): (1) **continuous monitoring of ALL four surfaces** — fatal crashes, ANRs, performance traces/regressions, AND non-fatals (skipping any one is a PASS-bluff; non-fatals are the silent class — every catch/fallback/recovered-error path is a real degraded UX and MUST be triaged, not ignored because the app did not crash); (2) **systematic-debugging of each (reproduce-before-fix)** per §11.4.102 Iron Law + §11.4.115 RED-baseline-on-the-broken-artifact (the recorded stacktrace/ANR-trace/non-fatal-context IS the §11.4.5/.69 captured defect-present evidence) BEFORE the fix — a fix without a prior reproducing test is a §11.4.43/.123 violation; unreproducible ⇒ deep-research-before-declaring-untestable §11.4.123/.150; (3) **a fix/improvement for every confirmed issue** against its proven root cause (§11.4.9/.43/.108), “acknowledged in console” is not a fix, muting a recurring issue without a tracked rationale is a §11.4.6/.90 violation; (4) **validation+verification regression-test coverage per closed issue** — in the SAME commit as the fix register a permanent §11.4.135 regression guard (§11.4.115 polarity test: RED_MODE=1 captures the recorded defect on the pre-fix artifact, RED_MODE=0 standing GREEN guard asserting ABSENT) exercising the same user-reachable path the stacktrace identifies, with rock-solid captured evidence §11.4.5/.69/.107/.123 + a closure log citing the console issue id/URL + root-cause + fix commit + the validation/verification test paths; a Crashlytics issue marked resolved WITHOUT a falsifiable regression test is FORBIDDEN (the silent-recurrence vector, §11.4.138 operator-escape class applied to field telemetry); (5) **regular cadence** reusing the §11.4.47 five-trigger set (pre-build/pre-flash/pre-distribute/pre-tag blocking; daily + post-deployment-burn-in non-blocking) — a one-time sweep never repeated is a violation; (6) **no false results, no bluff** (§11.4/.1/.6 — “no new issues” requires the enumerated monitored-surfaces+window §11.4.118; “fixed” requires the RED→GREEN flip §11.4.115/.130; a guard whose §1.1 mutation does not FAIL it is a bluff gate). Honest boundary §11.4.6: processing every recorded issue breaks the silent-recurrence vector — it does NOT prove zero remaining field defects nor replace §11.4.108/.40; an issue is “closed” only when its guard is GREEN on a clean artifact, never when the console mark is flipped. Classification: universal (§11.4.17) — the consuming project supplies its Crashlytics handle, console-access credential (§11.4.10, never logged), severity table, and per-issue closure-log path per §11.4.35; the reference project-level instantiation is the Lava client’s §6.O (Crashlytics-Resolved Issue Coverage — per-issue validation test + Challenge test + .lava-ci-evidence/crashlytics-resolved/<date>-<slug>.md closure log) + §6.AC (Comprehensive Non-Fatal Telemetry — every catch/fallback records a non-fatal with triage context). Composes §11.4/.1/.5/.6/.9/.34/.40/.43/.47/.69/.90/.102/.107/.108/.115/.118/.123/.130/.135/.138/.150/§1.1. Propagation gate CM-COVENANT-114-152-PROPAGATION (literal 11.4.152) + recommended gate CM-CRASHLYTICS-ISSUE-FULLY-COVERED (every closed Crashlytics issue carries a systematic-debug root-cause record + a registered §11.4.135 regression guard + a closure-log entry citing the console issue id/URL + the validation/verification test paths; a console-resolved issue with no falsifiable

regression guard FAILs) + paired §1.1 meta-test mutation (gate-code = separate work item). **Canonical authority:** constitution submodule Constitution.md §11.4.152. Non-compliance is a release blocker. No escape hatch — no `--skip-crashlytics-monitoring`, `--monitor-crashes-only`, `--skip-non-fatals`, `--resolve-without-regression-test`, `--console-mark-is-fixed`, `--monitor-once`, `--mute-without-rationale` flag.

§11.4.153 — Comprehensive per-feature Status + Status_Summary document set with mandatory video-recording confirmation (User mandate, 2026-06-15). Every project MUST maintain, under docs/features/, a comprehensive feature Status document set (Status.md + its §11.4.56 Status_Summary.md companion) enumerating EVERY system component, EVERY client app/binary/surface (TUI / CLI / Web / desktop / mobile / API / gRPC / library / submodule / infrastructure), and EVERY feature — including features ported from any incorporated CLI-agent / submodule catalogue (§11.4.74) — with NO single feature left out. (1) **Total categorized coverage** — per-feature table organized Component→Category, reconciled against the actual codebase (a code-present feature missing from the table, or a row with no code, is a §11.4.6/§11.4.118 bluff). (2) **Per-feature fields** — Component / Feature / Category / Implementation / Wiring (genuinely end-user-reachable, not merely compiled, §11.4.108) / Real-use / Tests-coverage (four-layer §11.4.4(b)) / Validation (PASS / FAIL / SKIP / PENDING_FORENSICS / OPERATOR-BLOCKED per §11.4.45) / **Video-recording confirmation** (path to the real-use video or honest gap marker). (3) **Mandatory per-feature real-use video** — every user-visible “confirmed” claim backed by a recorded real-use video of a genuine end-user scenario (real prompts → real LLM/service responses → real results), NEVER a frozen/stale frame (§11.4.107), NEVER faked/mock/demo-loop/bluff response or LLM error passed off as success (§11.4.2/§11.4.5), stored at the project-declared recording path (§11.4.35); a confirmed row with no real video or a bluff video = §11.4 PASS-bluff; autonomous-infeasible ⇒ honest §11.4.3/§11.4.52 SKIP + tracked migration item, NEVER a faked confirmation. (4) **Video-analysis remediation loop** — every video analysed; any defect it surfaces triggers §11.4.102 systematic-debug → fix → §11.4.146 retest → re-record → clean GO per §11.4.134 before “confirmed” (a video exposing a broken feature is a §11.4.4 test-interrupt). (5) **Always-in-sync** — §11.4.45-class roster/corpus-backed, §11.4.106 docs_chain-bound + §11.4.86 drift-proof fingerprint (sha256 of sorted feature-key roster AND sorted video-artefact roster, NOT mtime), re-syncs out-of-the-box; stale = violation. (6) **Four-format export** — HTML + PDF + **DOCX** (this doc class ADDS DOCX to the §11.4.65 HTML+PDF set; other classes unchanged), in sync per §11.4.60. (7) Follows §11.4.44/.45/.56/.57/.59/.60. (8) **MP4 format REQUIRED.** All video confirmations MUST be in .mp4 format (H.264). Window-specific capture ONLY (§11.4.159(A)). Vision validation REQUIRED (§11.4.159(D)). .cast files are supplementary only. Honest boundary §11.4.6 — guarantees a complete video-confirmed always-synced ledger, NOT per-feature correctness (rests on each row’s §11.4.69/.107/.123 evidence) and NOT a §11.4.40 retest substitute. Classification: universal (§11.4.17). Composes §11.4.2/.5/.44/.45/.52/.56/.57/.59/.60/.65/.86/.102/.106/.107/.108/.118/.123/.134/.146 /§1.1. Propagation gate CM-COVENANT-114-153-PROPAGATION (literal 11.4.153) + recommended gates CM-FEATURE-STATUS-COMPLETE + CM-FEATURE-STATUS-VIDEO-CONFIRMED + paired §1.1 mutation.

Canonical authority: constitution submodule Constitution.md §11.4.153. Non-compliance is a release blocker. No escape hatch — no `--skip-feature-status`, `--feature-without-video`, `--frozen-video-OK`, `--bluff-response-video-OK`, `--skip-video-analysis`, `--feature-ledger-incomplete-OK`, `--no-docx-export`, `--allow-feature-ledger-drift` flag.

§11.4.154 — Window-scoped capture + fresh-corpus rotation for feature/QA recordings (User mandate, 2026-06-15). Verbatim: “recording you perform does record the window containing apps and services, not the whole desktop or monitor screen!” + “All old recording files MUST BE removed when new one starts!” Refines §11.4.2/.5/.107/.153 recording discipline with two capture-hygiene invariants. **(A) Window-scoped, NOT whole-screen** — every feature/QA video MUST capture ONLY the window/surface of the app/service under test (GUI window / CLI-TUI terminal pane / web tab-viewport / device-emulator-simulator frame), NEVER the whole desktop/monitor or unrelated windows; whole-desktop capture leaks operator-private content (§11.4.10/.83), dilutes the §11.4.107 liveness/freeze oracle, and breaks the §11.4.137 OCR/ROI content oracle. Target the window/region by stable identity (window id/title, device serial, browser context, tmux target) per §11.4.111 — never a fixed full-screen device index. Platform genuinely cannot capture below whole-screen ⇒ honest §11.4.3 SKIP + tracked migration item, never a whole-screen pass-off. **(B) Fresh-corpus rotation** — when a new recording run for a scope begins, the agent’s OWN prior in-scope stale recordings at the raw recording path MUST be removed FIRST so the live corpus reflects the current run (§11.4.107 not-stale + §11.4.86 roster-freshness). Honest boundary (§11.4.6 + §9.2): “remove old” = the agent’s own prior recordings for the SAME scope/project ONLY — NEVER another project’s/scope’s/operator-authored files; uncertain ⇒ surface, don’t delete (§11.4.122); committed docs/qa/<run-id>/ evidence (§11.4.83) is the durable record, NOT rotated. Classification: universal (§11.4.17). Composes §11.4.2/.5/.10/.83/.86/.107/.111/.122/.128/.137/.153/ §9.2/.6. Propagation gate CM-COVENANT-114-154-PROPAGATION (literal 11.4.154) + recommended gate CM-WINDOW-SCOPED-FRESH-CORPUS-RECORDING + paired §1.1 mutation.

Canonical authority: constitution submodule Constitution.md §11.4.154. Non-compliance is a release blocker. No escape hatch — no --whole-screen-capture-OK, --skip-window-scope, --keep-stale-recordings, --no-corpus-rotation, --full-desktop-recording flag. **(C) MP4 auto-conversion REQUIRED.** Any .cast file produced MUST be auto-converted to .mp4 via agg + ffmpeg immediately after capture. The .mp4 is the primary evidence; .cast is supplementary only.

§11.4.155 — Project-name-prefixed feature/QA recording filenames (User mandate, 2026-06-15). Verbatim: “All recorded videos MUST START with prefix: the PROJECT NAME (ALWAYS USE THE PROJECT NAME). Project name MUST be obtained according to the constitution’s own project-name resolution.” Every recorded video the project produces — every feature/QA real-use recording (§11.4.153), every window-scoped capture (§11.4.154), every always-on device recording (§11.4.128), every raw or curated artefact at the project-declared recording path (§11.4.35) + the committed docs/qa/<run-id>/ trail (§11.4.83) — MUST have a filename that STARTS WITH the PROJECT-NAME prefix, ALWAYS; an unprefixed recording is a §11.4.155 violation (a multi-project/scope corpus on one host per §11.4.128/.103 becomes un-greppable + un-attributable — the §11.4.151 identify-and-grep failure on the recording-corpus axis). **Prefix resolution (closed-set, deterministic — §11.4.6, IDENTICAL to §11.4.151):** (1) HELIX_RELEASE_PREFIX from .env (authoritative, git-ignored §11.4.30, documented in tracked .env.example §11.4.77) else (2) lowercased snake_case project-root dir name §11.4.29 — ALWAYS resolvable, zero operator input. SAME prefix for EVERY recording in a checkout so ls '<PREFIX>---'* enumerates the corpus; canonical form <PREFIX>---<feature-or-scope>---<run-id>.<ext> (--- delimits the prefix unambiguously); MUST equal the §11.4.151-resolved release-tag prefix (divergence is itself a §11.4.155 violation — one project, one

name). Honest boundary (§11.4.6): the prefix guarantees attribution + greppability, NOT content validity (still §11.4.107/.137/.153) and does NOT relax §11.4.154's window-scope/rotation (rotation removes the agent's OWN <PREFIX>---* only; foreign/operator files surfaced not deleted §11.4.122/§9.2). Classification: universal (§11.4.17). Composes §11.4.151/.128/.153/.154/.111/.83/.6/.29/.30/.35/.77/.86/§1.1. Propagation gate CM-COVENANT-114-155-PROPAGATION (literal 11.4.155) + recommended gate CM-RECORDING-PROJECT-NAME-PREFIX + paired §1.1 mutation.

Canonical authority: constitution submodule Constitution.md §11.4.155. Non-compliance is a release blocker. No escape hatch — no --no-recording-prefix, --recording-without-project-name, --unprefixed-recording, --prefix-optional-for-recording, --differing-recording-prefix flag.

§11.4.156 — All CI/CD automation (GitHub Actions / GitLab pipelines / equivalents) MUST be disabled (User mandate, 2026-06-15). Verbatim operator mandate: “Any GitHub actions or GitLab pipelines MUST BE disabled! Add this critical mandatory rule / mandatory constraint into the root constitution Submodule, commit and fetch all its changes to all upstreams and make sure we respect and follow this rule (we do apply it) ASAP!!!” Every repository this Constitution governs — main repo, this constitution submodule, every owned + nested submodule we author and push — MUST ship with ALL server-side CI/CD automation DISABLED: no push to any owned upstream may trigger a GitHub Actions run, GitLab pipeline, or equivalent (Jenkins/CircleCI/Travis/Drone/Woodpecker/Bitbucket/Azure, any on: push/schedule/workflow_dispatch). GENERALISES + makes ABSOLUTE the §11.4.75 Layer-5 posture (remote CI DISABLED, workflow preserved at a ...disabled-local-only non-.yml name a provider ignores) across ALL governed repos; enforcement migrates to the LOCAL §11.4.75 git-hook ritual + §11.4.40 pre-tag sweep, never a remote runner. ALL hold: **(A)** zero active .github/workflows/*.yml|yaml / .gitlab-ci.yml / .gitlab/** / equivalent at the ROOT of any governed repo/submodule (the only place a provider executes); **(B)** “disabled” = a push triggers ZERO runs — delete OR rename to a non-trigger name (§11.4.75 .disabled/.disabled-local-only); a live-on:+if:false workflow still queues runs, NOT compliant; **(C)** scope = repos we author+push — vendored/third-party nested configs below the root (AOSP external/**, prebuilts/**, vendored submodules) are INERT (a provider never runs a non-root config), OUT of scope, MUST NOT be mass-edited (§11.4.29 vendor-exempt); test = “does a push to OUR upstream trigger a run?” yes⇒disable, inert⇒document+leave (§11.4.6 verify-not-assume); **(D)** no new CI may be added (release blocker); **(E)** pre-push verify git ls-files | grep -E '^\.github/workflows/.*\.ya?ml\$|^\.gitlab-ci\.yaml\$' empty for authored repos, §11.4.109-class PreToolUse guard + gate enforce mechanically. Honest boundary (§11.4.6): file-level disabling stops FILE-triggered runs, NOT provider-side server settings (org-default required workflows, branch-protection required checks, provider scheduled exports) — the operator turns those off; the agent documents what it cannot reach, never claims unachieved completeness. Composes §11.4.75/.29/.6/.40/.42/.109/.113/§2.1. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-156-PROPAGATION (literal 11.4.156) + recommended gate CM-NO-ACTIVE-CI + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; add a root .github/workflows/x.yml to an authored repo → CM-NO-ACTIVE-CI FAILs). **Canonical authority:** constitution submodule Constitution.md §11.4.156. Non-compliance is a release blocker. No escape hatch — no --allow-ci, --enable-workflow, --keep-pipeline, --remote-ci-OK, --ci-exempt flag.

§11.4.157 — GEMINI.md maintained in lockstep with CLAUDE.md / AGENTS.md / QWEN.md (User mandate, 2026-06-15). Verbatim operator mandate: “Make sure with CLAUDE.md, AGENTS.md, QWEN.md we maintain GEMINI.md too! Add this mandatory fact / rule to root constitution Submodule we are inheriting / extending - CONSTITUTION.md, CLAUDE.md, QWEN.md, AGENTS.md, GEMINI.md and other related relevant files!” Forensic FACT (2026-06-15): when §11.4.156/§11.4.157 were authored, Constitution/CLAUDE/AGENTS/QWEN carried the family through §11.4.155 but GEMINI.md had silently drifted to §11.4.141 — 14 mandates (§11.4.142–155) never propagated — a §11.4 propagation-bluff (a Gemini-CLI agent reads a stale Constitution). GEMINI.md is a FIRST-CLASS governance context carrier EQUAL to CLAUDE.md/AGENTS.md/QWEN.md, never optional/best-effort. ALL hold: **(A)** five-carrier lockstep — no governance change is complete until GEMINI.md carries it alongside the other three mirrors; GEMINI.md is added to the §11.4.26 propagation + cross-reference set explicitly; **(B)** no silent drift — GEMINI.md lagging the other mirrors’ highest rule is a §11.4.157 violation (§11.4.65-class), back-fill required; **(C)** equal status — GEMINI.md restates the SAME literal 11.4.N anchors the propagation gates require (§11.4.35), fleet count INCLUDES GEMINI.md; **(D)** consumer projects’ own CLAUDE/AGENTS/QWEN/GEMINI bind too (§11.4.35). Honest boundary (§11.4.6): the §11.4.142–155 GEMINI.md back-fill is a tracked release-blocking remediation; claiming GEMINI.md “in sync” while the back-fill is incomplete is itself a §11.4.157 violation. Composes §11.4.26/.35/.17/.44/.65/.140/.156/§1.1. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-157-PROPAGATION (literal 11.4.157, GEMINI.md INCLUDED) + recommended gate CM-GEMINI-MD-LOCKSTEP + paired §1.1 meta-test mutation. **Canonical authority:** constitution submodule Constitution.md §11.4.157. Non-compliance is a release blocker. No escape hatch — no --skip-gemini-md, --gemini-optional, --gemini-lag-OK, --four-carrier-suffices flag.

§11.4.158 — Intensive all-feature/flow/edge-case video-recording + read-the-screen content-verification mandate (User mandate, 2026-06-16). Every project MUST be covered by intensive automated testing that exercises + RECORDS every feature/flow/use-case/edge-case (valid/invalid/boundary/concurrent/failure-injection §11.4.85), each recording showing the feature GENUINELY WORKING with REAL results (real prompts→real responses→real outputs §11.4.153) and NO false/simulated/stale/frozen result (§11.4.2/.5/.107) — a recording showing an error/bluff/non-working feature is a FINDING (→§11.4.153(4)/§11.4.4 fix→retest→re-record), never a confirmation. The testing System MUST ACTUALLY READ every shown log line / message / UI label / dialog / toast / status text and VERIFY it is a genuine working result (OCR/ROI §11.4.117/.137 + confidence floor for pixel surfaces; direct capture for terminal/log; §11.4.107(10) self-validated golden-good/golden-bad analyzer) — “a video was produced” is NOT evidence, “the System read the screen + confirmed a genuine result” is. HelixQA (§11.4.27) MUST drive this exercise→record→read→score pass, PASS only on a read-confirmed genuine result + captured artefact path (§11.4.69). Default recording save path = \$HOME/Downloads (host user’s home Downloads, resolved at runtime never hardcoded) unless a project declares an override per §11.4.35; §11.4.155 project-prefix + §11.4.154 window-scope/rotation + §11.4.128 git-ignored raw corpus + §11.4.83 curated docs/qa evidence apply. **Vision analysis MANDATORY for every recording** — after every recording, the agent MUST read the terminal output / video content and verify the feature actually works. A recording without a vision-confirmed verdict is a §11.4.158 violation. Composes §11.4.2/.5/.25/.27/.52/.69/.83/.85/.107/.108/.117/.118/.128/.137/.138/.153/.154/.155/§1.1. Classification: universal (§11.4.17). Propagation gate CM-

COVENANT-114-158-PROPAGATION (literal 11.4.158) + recommended gates CM-INTENSIVE-RECORDING-COVERAGE + CM-RECORDING-CONTENT-READ-VERIFIED + paired §1.1 mutation. **Canonical authority:** constitution submodule Constitution.md §11.4.158. Non-compliance is a release blocker. No escape hatch — no --skip-recording-coverage, --video-without-content-read, --happy-path-recording-suffices, --recording-path-anywhere, --unread-recording-OK, --skip-helixqa-read flag.

§11.4.159 — Mandatory window-specific video recording + vision validation mandate (User mandate, 2026-06-20). Every feature test, validation, verification, challenge, and QA session that produces video evidence **MUST** comply with **ALL** of the following: **(A) Window-specific recording ONLY** — every video **MUST** record **ONLY** the target application window (Terminal pane, TUI app, browser tab, emulator frame), **NEVER** the whole desktop/monitor screen; use macOS screenshot -l<window_id>, Linux xdotool+ffmpeg, or equivalent; whole-desktop capture leaks operator-private content (§11.4.10), dilutes the §11.4.107 liveness oracle, and breaks §11.4.137 OCR/ROI; platform genuinely cannot capture below whole-screen => honest §11.4.3 SKIP + tracked migration item. **(B) MP4 format REQUIRED** — .mp4 (H.264, movflags +faststart, pix_fmt yuv420p); .cast files are supplementary only; auto-convert via agg+ffmpeg per §11.4.154(C). **(C) Project-name prefix REQUIRED** — filename **MUST** start with project name in snake_case from HELIX_RELEASE_PREFIX in .env (§11.4.151) or lowercased project root dir name (§11.4.29); format <project_name>--<feature>--<timestamp>.mp4; unprefixed = violation. **(D) Mandatory vision validation** — after **EVERY** recording, the agent **MUST** read the terminal output / video content and verify: (i) feature **ACTUALLY WORKS**, (ii) LLM responses are **REAL**, (iii) all tests show **PASS**, (iv) no “TODO implement” / “simulate” / “for now” patterns, (v) output demonstrates end-user working feature; **MUST** produce a verdict (**PASS**/**FAIL**) with evidence path (§11.4.69). **(E) Terminal window cleanup** — after each recording, dismiss/close **ONLY** the Terminal window used for that recording (window-specific close via osascript/xdotool); **MUST NOT** close windows belonging to other processes. **(F) Real results ONLY** — every recording **MUST** show **REAL** working features; errors/empty output/simulated responses trigger fix→retest→re-record. **(G) Re-runnable evidence** — the command shown **MUST** be re-runnable to produce the same results. **(H) Fresh-corpus rotation** — remove agent’s own prior in-scope stale recordings **FIRST** (§11.4.154). **(I) Content verification MANDATORY — not duration-based** — the value of a recording is **NOT** its duration but its **CONTENT**; a recording **MUST** demonstrate the **ACTUAL** feature being used with **REAL** results; before accepting **ANY** recording, the agent **MUST** verify: expected output patterns **ARE** present (e.g., test **PASS** lines, API response data, feature-specific output), the feature **ACTUALLY WORKS** as demonstrated (not just “something ran”), LLM responses are **REAL** content (not simulated, not placeholder, not empty), every claim of “working” is backed by visible evidence; a 5-second recording showing a feature working correctly is **MORE** valuable than a 60-second recording of empty terminal; duration is **NOT** a proxy for quality. **(J) Expected-content specification REQUIRED** — before recording, the agent **MUST** specify what content **SHOULD** appear (expected patterns, expected test results, expected API responses); after recording, the agent **MUST** verify these patterns **ARE** present; if expected content is **MISSING**, the recording is **REJECTED** regardless of duration. **(K) Content-verification recording workflow** — mandated workflow: (1) **SPECIFY** expected content patterns, (2) **RECORD** the feature execution, (3) **EXTRACT** all text from the recording, (4) **VERIFY** expected patterns are present, (5) **CHECK** for simulated/placeholder content, (6) **ACCEPT** only if **ALL** patterns found **AND** zero bluffs

detected, (7) REJECT and re-record if ANY pattern missing or bluff detected. **(L) Root cause analysis REQUIRED for rejected recordings** — when a recording is rejected (missing expected content, bluff detected, or empty capture), the agent MUST investigate WHY before re-recording per §11.4.102; determine the root cause (timing issue, wrong command, tool failure, etc.) and fix it; simply re-recording without understanding WHY the first attempt failed is a §11.4.159 violation. **(M) Real-time monitoring RECOMMENDED** — for complex features, use real-time monitoring that analyzes output DURING recording (not after); this catches issues immediately and allows corrective action before the recording completes. Classification: universal (§11.4.17). Composes §11.4.2/.3/.5/.10/.29/.69/.83/.107/.111/.128/.137/.151/.153/.154/.155/.158/§1.1. Propagation gate CM-COVENANT-114-159-PROPAGATION (literal 11.4.159) + recommended gate CM-WINDOW-VIDEO-VALIDATED + paired §1.1 mutation. **Canonical authority:** constitution submodule Constitution.md §11.4.159. Non-compliance is a release blocker. No escape hatch — no --whole-screen-ok, --cast-only, --skip-vision-validation, --no-cleanup, --simulated-recording-ok, --unprefixed-recording flag.

§11.4.160 — Vision-verified recording + HelixQA bridge mandate (User mandate, 2026-06-21). Compact summary: every video recording for feature/QA evidence MUST be processed through a vision/OCR pipeline that reads on-screen content and confirms expected results BEFORE acceptance; the recording system MUST provide a bridge feeding captured frames to HelixQA's test infrastructure (or equivalent) for automated read-the-screen verification against SPECIFY-phase expected patterns (§11.4.159(J)); the bridge MUST capture frames at ≤5s intervals, run OCR/vision analysis with a self-validated golden-good/golden-bad analyzer (§11.4.107(10)), compare extracted text against specified patterns, produce a per-frame PASS/FAIL verdict with an evidence path to the frame, and surface failures immediately so the recording can be re-done per §11.4.159(L). The capture interval, OCR confidence floor, and pattern-matching thresholds are project-configured per §11.4.35, calibrated on the project's own fixtures (§11.4.6). Honest boundary: vision verification confirms on-screen content — does NOT replace §11.4.5 quality analysis nor §11.4.108 runtime-signature verification; FLAG_SECURE surfaces use the §11.4.117 proxy oracle per §11.4.112. Classification: universal (§11.4.17). Composes §11.4.5/.27/.69/.107(10)/.112/.117/.153/.158/.159(J)/§1.1. Propagation gate CM-COVENANT-114-160-PROPAGATION + recommended gate CM-VISION-VERIFIED-RECORDING-BRIDGE + paired §1.1 mutation.

§11.4.161 — Rootless container runtime mandate (User mandate, 2026-06-21). Compact summary: every project MUST use Podman in rootless mode (or equivalent rootless container runtime) for ALL containerized workloads — Docker in rootful mode, sudo, or any escalation to root for container management is STRICTLY FORBIDDEN unless the target platform has no rootless option AND that constraint is documented per §11.4.112; the vasic-digital/containers Submodule (§11.4.76) MUST be used as the sole container orchestration layer — no ad-hoc docker/podman commands outside the Submodule's pkg/boot/pkg/compose/pkg/health primitives; if a missing capability forces a raw command, the containers Submodule MUST be extended upstream per §11.4.74 rather than worked around; all container-related integration tests MUST boot infra on-demand via the Submodule (the on-demand-infra invariant). Honest boundary: rootless execution eliminates the container-to-root privilege-escalation vector — it does NOT replace §11.4.10 credentials-handling (credentials in containers still require leak audits) nor §12.3 container hygiene (memory limits/OOM policies/restart backoff still apply). Classification: universal (§11.4.17). Composes

§11.4.76/.74/.10/.112/§12.3/.6. Propagation gate CM-COVENANT-114-161-PROPAGATION + recommended gate CM-ROOTLESS-CONTAINER-RUNTIME + paired §1.1 mutation.

§11.4.162 — OpenDesign UI design system mandate (User mandate, 2026-06-21). Compact summary: every project producing user-facing interfaces (web, desktop, mobile, TUI) MUST use OpenDesign (<https://github.com/nexu-io/open-design>) as the mandatory UI design-and-refinement system — NOT ad-hoc CSS or one-off design tools; install as a project dependency and use its design tokens/themes system for: (a) all color palette definitions supporting BOTH light and dark themes from project brand assets (§11.4.35), (b) typography scale and spacing, (c) component-level design tokens; if a desired UI pattern is not supported, extend OpenDesign upstream per §11.4.74 (extend-don't-reimplement); every UI component MUST ship light + dark theme variants; elements MUST NOT overlap, fonts MUST NOT collide, labels MUST NOT overlay labels — any layout regression is a §11.4.162 violation; all UI changes MUST be covered by the project's standard test types including visual regression tests (before/after screenshots with per-pixel or perceptual-diff PASS/FAIL). Honest boundary: OpenDesign governs design tokens and themes — it does NOT replace functional testing (§11.4.27), WCAG accessibility assertions (§11.4.107), nor the §11.4.48/§11.4.49 UI-driven and dual-approach testing methodology. Classification: universal (§11.4.17). Composes §11.4.74/.25/.27/.4(b)/.107/.48/.49/.35/.69/§1.1. Propagation gate CM-COVENANT-114-162-PROPAGATION + recommended gate CM-OPENDESIGN-UI-SYSTEM + paired §1.1 mutation.

§11.4.163 — Universal Media Validation & Verification Mandate (User mandate, 2026-06-21). Compact summary: every recorded artifact (video/audio/screenshots/asciicast/text) MUST pass a MEDIA VALIDATION pipeline before acceptance — OCR for video/screenshots (§11.4.117/.107(12)), transcription for audio, text parsing; compare extracted content against SPECIFY-phase expected patterns (§11.4.159(J)); self-validated analyzer with golden-good/golden-bad fixture pair (§11.4.107(10)); produce structured verdict (PASS/FAIL + evidence path + matched/unmatched patterns + pinpoint data on FAIL — which frame/line/timestamp, expected vs actual); post-recording AND real-time trigger (§11.4.159(M)); paired §1.1 mutation ensures golden-bad fixture produces FAIL. Honest boundary: confirms artifact content matches patterns — does NOT replace §11.4.5 quality, §11.4.107 liveness, nor §11.4.108 runtime-signature. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-163-PROPAGATION + recommended gate CM-MEDIA-VALIDATION-PIPELINE + paired §1.1 mutation.

§11.4.164 — Universal Constitution Auto-Propagation & Hook System (User mandate, 2026-06-21). Compact summary: every fetch+pull of the constitution submodule MUST trigger constitution/scripts/post_update_hook.sh (inherited §11.4.28) that: detects changed files (governance MDs + scripts/hooks/skills/mcp/plugins); registers new/modified skills via consumer's scripts/register_skills.sh; registers MCP servers via consumer's scripts/register_mcp.sh; installs hooks into .git/hooks/; validates scripts syntax per §11.4.67; emits summary log; failures logged per §11.4.6; tested by §1.1 mutation (add skill → hook detects it). Consumer MUST invoke after every pull. Honest boundary: installs components — does NOT guarantee runtime acceptance nor replace §11.4.32 sweep. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-164-PROPAGATION + recommended gate CM-CONSTITUTION-AUTO-PROPAGATION + paired §1.1 mutation.

§11.4.165 — Universal Independent Verification Agent Mandate (User mandate, 2026-06-21). Compact summary: every code/media/docs/config output MUST pass INDEPENDENT verifier (§11.4.70/.20, structurally separate from author, §11.4.92 PRECEDES never satisfies), iterating to GO per §11.4.134. CODE: §11.4.142 review + build+test + paired §1.1 mutations + §11.4.108 runtime-signature. MEDIA: §11.4.163 pipeline + genuine-content check (§11.4.159(I)) + format check (§11.4.159(B)). DOCS: exports current §11.4.65 + revision header §11.4.44. CONFIG: syntax + schema + §11.4.10 leak check. Structured findings per class; verifier self-validated by §1.1 mutation. Honest boundary: source/artifact integrity — does NOT replace §11.4.108 runtime-signature nor §11.4.40 retest. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-165-PROPAGATION + recommended gate CM-INDEPENDENT-VERIFICATION-AGENT + paired §1.1 mutation.