

Helix Universal Constitution

Field	Value
Revision	65
Created	2026-05-14
Last modified	2026-08-18T18:00:00Z
Status	active

Status summary	Amendment of 2026-08-15 — MONOLITHIC ROUND: 22 NEW universal anchors (§11.4.239-254 AI-curriculum modules 27-35 harvest + §11.4.257-262 enterprise-quality 9-point mandate) + 5 targeted extensions (§11.4.212 EXT clauses (E)/(F) README intro+illustration + linkage-completeness ratchet; §11.4 covenant preamble EXT machine-created evidence at every gate binding; §11.4.108 EXT clause (7) machine-created evidence at every layer; §11.4.224 EXT (B) 7-canonical-type coverage floor + 100% aim; §11.4.27 EXT (B) 7-canonical-type breadth enumeration). §11.4.239 critical-invariant DoD money/safety/availability/integrity failure-path scenarios as gates; §11.4.240 producer≠verifier scope separation via least-privilege-not-instruction; §11.4.241 illegal-state-unrepresentability preference ladder (types → API-shape → lint → property-test → runtime-assert); §11.4.242 bisection-not-blame regression cause-finding + failure-mode catalogue; §11.4.243 characterization/golden-master safety net before behaviour-changing legacy modification; §11.4.244 cross-boundary contract tests + can-i-deploy gate merged; §11.4.245 oracle-problem-first test authoring with named strategy + independence; §11.4.246 reproducible + hermetic builds + SLSA Build Level 2 fleet-wide minimum + supply-chain integrity chain; §11.4.247 composite-output layer-move must move every layer with completeness check + compatibility window; §11.4.248 flaky-test quarantine + [PROTECTED-SPEC: ATM-NNN] tag + CODEOWNERS required reviewer; §11.4.249 producer≠oracle≠gate≠verifier four-role separation + flight recorder; §11.4.250 heuristic-tower signals a primitive defect requiring systematic-debug not layer N+1; §11.4.251 byte-identical-fork prohibition + role-as-data-pack extraction; §11.4.252 fail-closed-on-dangerous-combination for ≥2 mutation/untrusted-input/credential/side-effect/exec/irreversible capabilities; §11.4.253 idempotency under retry with DB-level durable uniqueness guard; §11.4.254 boot-time invariant assertion + capability matrix; §11.4.257 comprehensive user-manual + guide + FAQ coverage per component/service/feature; §11.4.258 architectural + data-flow + state-machine + sequence-diagram coverage embedded not orphan; §11.4.259 README quality-status badge row with closed green → amber → red vocabulary + machine-derived + production-readiness gauge; §11.4.260 cutting-edge enterprise quality + production-deployment readiness 10-invariant tracker; §11.4.261 zero-shortcomings/gaps/weak-spots/danger-zones invariant + mechanical audit ratchet monotone-decreasing; §11.4.262 machine-created evidence at every gate binding covenant preamble universally. SUBSTRATE-GAP DISCLOSURE per §11.4.231(F)(b) honest boundary: the §11.4.209 review substrate for this round ran on the opus-class
----------------	---

parent-inherited fork substrate (the Claude Code Agent tool exposes a `model` parameter but NO `effort` parameter, so subagent effort is not presently settable via that path — the label `<effort>` field degrades honestly to `? )` — SUBSTANTIVE on structural/traceability/composes/anti-bluff/classification/collision/cross-anchor/no-escape-hatch axes with clean GO on round 2 (subagent a6b89ff0, all IMPORTANT-1/2/3 + MINOR-1/2/3/4 findings resolved, zero new BLOCKING or IMPORTANT introduced); a re-review at the true §11.4.209 Fable- `xhigh` substrate via the Workflow-tool `agent()` path REMAINS RECOMMENDED (not blocking) before the anchors are considered fully substrate-verified — tracked as an owed §11.4.197 upgrade item, never a claim of shipped Fable-`xhigh` review. Classification: universal (§11.4.17). Propagation gates CM-COVENANT-114-239 through CM-COVENANT-114-254 + CM-COVENANT-114-257 through CM-COVENANT-114-262 (14 anchor-block-start propagation gates + 6 more) + recommended mechanism gates enumerated per anchor + paired §1.1 mutations (gate-code = separate work item, NOT claimed shipped §11.4.6/ §11.4.227). §11.4.239-254 + §11.4.257-262 amendment region byte-identical across the 4 mirrors CLAUDE.md + AGENTS.md + QWEN.md + GEMINI.md per §11.4.157 lockstep; Constitution.md carries the canonical ### §11.4.NNN heading variant per §11.4.227(B) canonical-vs-mirror layer variance legitimate. All 22 new anchors exactly-once block-start per file per §11.4.227(B). Generated .html/.pdf/.docx twins NOT regenerated (exporter still not a runnable in-repo script — honestly STALE per §11.4.106(E)). Prior round: Amendment of 2026-08-05 — NEW universal anchor §11.4.237 (operator mandate 2026-08-05): every translation / translated artifact (UI strings, pages, documents, PDFs, any localized content) MUST undergo a HEAVY, EXHAUSTIVE, INDEPENDENT context-and-spirit-aware review by a reviewer STRUCTURALLY SEPARATE from the translator (never self-review — §11.4.92 precedes, §11.4.142 never satisfies); the reviewer holds + applies (a) full CONTEXT of the source material (original ideas / subject-domain / purpose / audience), (b) the SPIRIT / idiom / register / tradition / cultural-conventions of the TARGET language so the result reads NATIVE + natural + maximally polished, never literal/word-for-word/calqued, and (c) correct TECHNICAL vocabulary + VERBATIM preservation of NON-TRANSLATABLE terms (product/brand/proper names, technology/library/API names, code identifiers); goal = maximally slick, polished, professional localization — accuracy NECESSARY, never SUFFICIENT; anti-bluff — captured machine-readable POSITIVE evidence (verdict + per-dimension rationale across {fidelity, fluency/naturalness, spirit/register, terminology,

completeness, non-translatable-preservation}, feature_class=translation_review, §11.4.5/§11.4.69) and a FAIL on any dimension loops back to re-translation/polish until a zero-finding clean verdict (§11.4.134). REFINES §11.4.140 (canonical translation pipeline) / §11.4.141 (independent per-language translation review — from PASS/FAIL accuracy to context+spirit native polish; the KNOWN §11.4.54 anchor-number collision on 11.4.140/11.4.141 cited strictly in those senses — re-mint RESOLVED 2026-08-15 (operator-owned decision executed: HelixTranslate canonical translation pipeline → §11.4.255, independent per-language translation review → §11.4.256; Constitution.md 5ed8c80 + mirrors CLAUDE 6dff307 / AGENTS f497412 / QWEN 9e45dec / GEMINI e5f2891), NO new collision introduced). Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-237-PROPAGATION + recommended gate CM-TRANSLATION-CONTEXT-SPIRIT-REVIEW + paired §1.1 mutation (accuracy-only check / self-reviewed-by-own-translator / golden-bad mistranslated-idiom or "translated" non-translatable product name → gate FAILs; gate-code = separate work item, NOT claimed shipped §11.4.6/ §11.4.227). §11.4.237 amendment region byte-identical across the 4 mirrors CLAUDE.md + AGENTS.md + QWEN.md + GEMINI.md (full-block md5 05d4e829de58414ca8f7ce920b606173, §11.4.157 lockstep); Constitution.md carries the canonical ####-heading variant (identical body md5 8f518d2f25d8c61a94b114c32c2db308 — §11.4.227(B) canonical-vs-mirror layer variance legitimate, content-hash equality binds within the lockstep set); §11.4.237 exactly-once block-start per file (§11.4.227(B)), no anchor-number collision. Generated .html/.pdf/.docx twins NOT regenerated (exporter still not a runnable in-repo script — honestly STALE per §11.4.106(E)). Prior round: Amendment of 2026-07-26 — NEW universal anchor §11.4.234 (operator mandate 2026-07-26, verbatim: "We MUST have dedicated script to run hook validations, and commit and push mechanism ALWAYS unblocked!"): hook validations run as an EXPLICIT stage of a dedicated, idempotent commit/push script (consumer-bound path as DATA per §11.4.35 — Lava binding: scripts/commit-push-all.sh); automatic git hooks wired via core.hooksPath MUST NOT gate routine commit/push (disconnect = hooksPath unset, hook file preserved UNMODIFIED, --no-verify never the routine path); NO gate is lost (every disconnected check remains executed at a named script stage or release/tag gate, silently dropping one is a §11.4 PASS-bluff at the process layer); the commit/push mechanism is ALWAYS unblocked (failing validation = clear per-check report + documented remediation path, long multi-minute gates separable

Field	Value
	from cheap checks and skippable ONLY via an explicit recorded deferral flag — Lava binding: LAVA_SYNC_SKIP_CI=1 — never an opaque hung/rejected push). Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-234-PROPAGATION + recommended mechanism gates CM-DEDICATED-HOOK-VALIDATION-SCRIPT / CM-HOOKS-NEVER-BLOCK-PUSH / CM-NO-GATE-LOST-ON-HOOK-DISCONNECT / CM-COMMIT-PUSH-ALWAYS-UNBLOCKED + paired §1.1 mutations (gate-code = separate work item, NOT claimed shipped §11.4.6/§11.4.227). Amendment region byte-identical across all 5 governance files (md5-verified); §11.4.234 exactly-once block-start per file (§11.4.227(B)), no anchor-number collision. CLAUDE.md + AGENTS.md + QWEN.md + GEMINI.md mirrors updated in lockstep (§11.4.157). Generated .html/.pdf/.docx twins NOT regenerated (exporter still not a runnable in-repo script — honestly STALE per §11.4.106(E)). Prior round: Amendment of 2026-07-26 — the EFFORT-tiering + 5-field-label batch (operator IMPORTANT mandate 2026-07-26, verbatim: "Besides dynamic determining proper models to use, we MUST decide the effort as well and switch between the effort settings same as we do now for chosen model! Labeling MUST BE extended for this details as well. Example: (T1/main - claude4 - opus - high)."). §11.4.182 EXTENSION amended (in-place, no re-mint): the work-stream label form extends from (T<N>/<branch> - <alias> - <model>) to (T<N>/<branch> - <alias> - <model> - <effort>) , effort ∈ {low
Is-sues	none
Is-sues summary	ToC is currently stale (missing entries for §11.4.71/.72/.73/.74/.75/.76) — separate clean-up commit needed per §11.4.61.
Fixed	§11.4.108 added; §11.4.73, §11.4.74, §11.4.76; QWEN.md created
Fixed summary	Operator-mandated additions 2026-05-20 propagated into Constitution.md (this commit); CLAUDE.md + AGENTS.md + new QWEN.md mirrors land in the same commit. §11.4.76 renumber-from-§11.4.75 documented inline (collision with concurrent Mechanical Enforcement landing).

Field	Value
Continuation	(1) ToC regeneration commit (cover §11.4.71..§11.4.76). (2) Submodule-catalogue inventory expansion in CLAUDE.md / AGENTS.md / QWEN.md per the 2026-05-20 follow-up user mandate — per-Submodule purpose lines for every repo in <code>vasic-digital</code> + <code>HelixDevelopment</code> .

Table of contents

- 1. Test coverage is mandatory for every change
 - 1.1 False-positive immunity is an invariant
- 2. Commit and push mechanics — single entryptpoint, locked
 - 2.1 Multi-upstream push is the norm
- 3. Submodule changes propagate through submodule commits first
- 4. Every tag on the main repo MUST be mirrored on every owned submodule
- 5. Changelog discipline and multi-format export
- 6. Documentation up to the nano-details
- 7. Making false-success results literally impossible
 - §7.1 NO BLUFF — positive-evidence-only validation
- §8. Bleeding-edge ultra-perfection quality bar
- §9. Absolute codebase and data safety — zero risk, zero loss
 - §9.1 Mandatory safety protocol for destructive operations
 - §9.2 Force-push requires explicit user authorization every time
 - §9.3 Hardlinked backup is the standard — there is no excuse
 - §9.4 Commit-message audit trail for history rewrites
- §10. Enforcement
- §11. End-user quality covenant
 - §11.4 End-user quality guarantee — forensic anchor (User mandate, 2026-04-28)
 - §11.4.1 — FAIL-bluffs are equally forbidden

- §11.4.2 — Recorded-evidence requirement
- §11.4.3 — Per-environment-topology test dispatch
- §11.4.4 — Test-interrupt-on-discovery + retest-from-clean-baseline
- §11.4.5 — Captured-evidence quality analysis
- §11.4.6 — No-guessing mandate
- §11.4.7 — Demotion-evidence rule
- §11.4.8 — Deep-web-research-before-implementation
- §11.4.9 — Batch-source-fixes-before-rebuild
- §11.4.10 — Credentials-handling mandate
- §11.4.10.A — Pre-store credential leak audit (User mandate, 2026-05-17)
- §11.4.11 — File-layout discipline
- §11.4.12 — Auto-generated docs sync mandate
- §11.4.13 — Out-of-band sink-side captured-evidence
- §11.4.14 — Test playback cleanup mandate
- §11.4.15 — Item-status tracking mandate
- §11.4.16 — Item-type tracking mandate
- §11.4.17 — Universal-vs-project classification of new rules (User mandate, 2026-05-14)
- §11.4.20 — Subagent-driven-by-default mandate (User mandate, 2026-05-14)
- §11.4.18 — Script documentation mandate (User mandate, 2026-05-14)
- §11.4.19 — Fixed-document column-alignment mandate (User mandate, 2026-05-14)
- §11.4.21 — Operator-blocked status + self-resolution exhaustion mandate (User mandate, 2026-05-14)
- §11.4.22 — Document-sync commit discipline (User mandate, 2026-05-14)
- §11.4.23 — Visual-cue & grouping mandate for Issues docs (User mandate, 2026-05-14)
- §11.4.24 — Build-resource stats tracking mandate (User mandate, 2026-05-14)
- §11.4.25 — Full-Automation-Coverage Mandate (User mandate, 2026-05-15)

- §11.4.26 — Constitution-Submodule Update Workflow Mandate (User mandate, 2026-05-15)
- §11.4.31 — Submodule-Dependency-Manifest Mandate (User mandate, 2026-05-15)
- §11.4.32 — Post-Constitution-Pull Validation Mandate (User mandate, 2026-05-15)
- §11.4.30 — .gitignore + No-Versioned-Build-Artifacts Mandate (User mandate, 2026-05-15)
- §11.4.29 — Lowercase-Snake_Case-Naming Mandate (User mandate, 2026-05-15)
- §11.4.28 — Submodules-As-Equal-Codebase + Decoupling + Dependency-Layout Mandate (User mandate, 2026-05-15)
- §11.4.27 — No-Fakes-Beyond-Unit-Tests + 100%-Test-Type-Coverage Mandate (User mandate, 2026-05-15)
- §11.4.33 — Type-aware closure-status vocabulary (User mandate, 2026-05-15)
- §11.4.34 — Reopened-source attribution mandate (User mandate, 2026-05-15)
- §11.4.35 — Canonical-root inheritance clarity (User mandate, 2026-05-15)
- §11.4.36 — Mandatory install_upstreams on clone/add Mandate (User mandate, 2026-05-15)
- §11.4.37 — Fetch-before-edit mandate (User mandate, 2026-05-15)
- §11.4.38 — Installable-Asset Evidence Mandate (User mandate, 2026-05-17)
- §11.4.39 — Per-Feature On-Device End-User Validation Mandate (iter-76, 2026-05-17)
- §11.4.40 — Full-suite retest before release tag mandate (User mandate, 2026-05-17)
- §11.4.41 — Pre-Force-Push Merge-First Mandate (User mandate, 2026-05-17)
- §11.4.42 — Iteration-discipline mandate (User mandate, 2026-05-18)
- §11.4.43 — TDD-Fix-Discipline Mandate (User mandate, 2026-05-18)

- §11.4.44 — Document Revision Header Mandate (User mandate, 2026-05-18)
- §11.4.45 — Integration-Status-Doc Maintenance Mandate (User mandate, 2026-05-18)
- §11.4.46 — Validate-recent-work-before-post-flash-tests mandate (User mandate, 2026-05-18)
- §11.4.47 — Firebase Data Review Mandate (User mandate, 2026-05-18)
- §11.4.48 — UI-Driven Video Testing Mandate (User mandate, 2026-05-18)
- §11.4.49 — Dual-Approach Testing Mandate (User mandate, 2026-05-18)
- §11.4.50 — Deterministic Consistency Mandate (User mandate, 2026-05-18)
- §11.4.51 — Live-ADB-First Maximization Mandate (User mandate, 2026-05-18)
- §11.4.52 — Autonomous-Validation Mandate (User mandate, 2026-05-18)
- §11.4.53 — Fixed_Summary parity mandate (User mandate, 2026-05-18)
- §11.4.54 — ATM-NNN ticket identifier mandate (User mandate, 2026-05-19)
- §11.4.55 — Reopens-history tracking + per-item Reopens.md doc (User mandate, 2026-05-19)
- §11.4.56 — Status_Summary parity + two-audience format (User mandate, 2026-05-19)
- §11.4.57 — README.md doc-link section + revision metadata (User mandate, 2026-05-19)
- §11.4.58 — Parallel-development methodology (User mandate, 2026-05-19)
- §11.4.59 — README always-sync mandate (User mandate, 2026-05-19)
- §11.4.60 — Documentation always-sync composite covenant (User mandate, 2026-05-19)
- §11.4.61 — Mandatory Markdown metadata table + structured-doc ToC (User mandate, 2026-05-19)
- §11.4.63 — Workable-items procedure docs as single source of truth (User mandate, 2026-05-19)
- §11.4.65 — Universal Markdown export mandate (User mandate, 2026-05-19)

- §11.4.66 — Blocker-resolution interactive-clarification mandate (User mandate, 2026-05-19)
- §12. Host-session safety — directly OR indirectly signing the user out is FORBIDDEN
 - §12.1 Forbidden operations — directly OR indirectly
 - §12.2 Required safeguards
 - §12.3 Container hygiene
 - §12.6 Memory-Budget Ceiling — 60% MAXIMUM
 - §12.10 Continuation document — sacred invariant
- Appendix A — Mutation testing — academic and industrial foundations
- Appendix B — Recursive inheritance & path-independence
- Appendix C — Multi-upstream push

This document defines mandatory, non-negotiable rules for every project that includes this constitution submodule. Every AI agent and every human contributor MUST comply. Violations are blockers.

Project-specific constitutions extend this document. When a project Constitution (e.g. `docs/guides/PROJECT_CONSTITUTION.md`) and this universal Constitution differ, the project Constitution may extend or tighten — but NEVER weaken — the universal rules below.

Last revision: 2026-05-14 (v1.0 — initial extraction from a consuming project's Android app 1.1.5-dev)

This Constitution is the universal source-of-truth for the engineering practices shared across all Helix projects (and any project that opts in by adding this submodule). It is intentionally agnostic about specific hardware, operating systems, languages, vendors, and product domains. Anything in this document that mentions a specific product, hardware SKU, library version, or vendor MUST be moved to the consuming project's own Constitution.

. Test coverage is mandatory for every change

Every code, configuration, documentation, or infrastructure change MUST ship with tests that prove:

Invariant	Mechanism
The change is present in source	pre-build / pre-merge gate that scans the source tree for the expected pattern
The change survives compilation / packaging	post-build gate that inspects the produced artifact (binary, archive, image) and verifies the change actually shipped
The change behaves correctly at runtime	runtime / integration / on-device test that exercises the user-visible behaviour
The gate itself is not bluffing	meta-test (mutation) entry that mutates the asserted condition and proves the gate catches the break

A change without all four layers of coverage is not ready to merge.

. False-positive immunity is an invariant

A test that always returns PASS because its regex never matches, its input path is wrong, its assertion target does not exist, or its comparison is tautological is **worse than no test**. Every new gate MUST be paired with a mutation entry in the project's meta-test harness that:

1. Temporarily breaks the assertion (sed out a line, rename a symbol, etc.).
2. Re-runs the gate.
3. Asserts the gate now reports FAIL.
4. Restores the original.

If the mutation round does not turn PASS → FAIL, the gate is a sham and must be rewritten.

This is the **single most important rule in this Constitution**. Every subsequent §11.4.x clause is downstream of it.

. Commit and push mechanics – single entrypt, locked

All commit and push work uses the project's official multi-remote commit wrapper (e.g. `scripts/commit_all.sh`) as the single entrypt. Direct `git commit` / `git push` / `git add` on the main repo is prohibited in normal workflow.

Exception: tag creation + tag-ref push is done via explicit `git tag -a` and `git push <remote> <tag>`.

The commit and push wrappers MUST hold an advisory `flock` so two invocations cannot race against each other. Lock files self-clean on process exit via `trap`.

When a contributor sees the "another commit/push wrapper is already running" error, they MUST NOT `rm -f` the lock unless they have just killed the owning process.

. Multi-upstream push is the norm

Every project hosted on multiple Git providers (GitHub primary, GitLab / GitFlic / GitVerse / Bitbucket / Gitea mirrors) MUST push every commit to ALL configured upstream remotes. A commit that lives on only one remote is a future operational risk: when that one provider has an outage, the project is unreachable.

The constitution submodule itself ships an `install_upstreams.sh` helper (and an `Upstreams/` directory with one declaration script per remote). Running it from the submodule root configures all the remotes locally. Consuming projects SHOULD do the same for their own multi- remote topology.

. Submodule changes propagate through submodule commits first

When a change lands in any submodule of the project:

1. **Commit inside the submodule first** (the submodule's own `git add + git commit`), since the parent project's commit wrapper does not commit submodule source — it only captures the updated submodule pointer.
2. **Push the submodule commit** to **all** remotes of that submodule.
3. **Then** run the parent project's commit wrapper to capture the updated submodule pointer in main and push main.

Skipping step 1 produces parent commits / tags that point at old submodule HEADs without the actual source.

4

. Every tag on the main repo MUST be mirrored on every owned submodule

When a version tag is created on the main repo at a specific commit, the same tag MUST be created on **every owned submodule** at that submodule's currently-pointed-to HEAD, and pushed to **every remote** of every owned repo. Third-party submodules (libraries not under the project's control) are excluded. The consuming project's Constitution declares its owned-submodule set explicitly.

5

. Changelog discipline and multi-format export

Every tagged release MUST ship:

1. A new `docs/changelogs/<tag>.md` entry describing user-visible changes, developer-visible changes, and known caveats.

2. Exports of that entry to `docs/changelogs/<tag>.{html,json,txt}` so external tooling (release portals, ticketing systems, package indices) can consume without needing a Markdown parser.
3. An update to a cumulative `docs/changelogs/CHANGELOG.md` (reverse chronological).

The project SHOULD provide a script (`scripts/testing/export_changelog.sh` `<tag>` or equivalent) that wraps `pandoc + jq + plain-text` Markdown strip.
6

. Documentation up to the nano-details

Every new feature, fix, or infrastructure change MUST update:

1. The project's CLAUDE.md / AGENTS.md Applied Fixes table (one row, one commit, one release tag).
2. `docs/guides/` for any user-visible or developer-reachable behaviour change.
3. Architecture diagrams, flowcharts, and any reference docs that touch the changed subsystem.
4. The per-version changelog (§5).

Documentation drift after a fix is itself a Constitution violation: undocumented behaviour change is the same defect class as un-tested behaviour change.
7

. Making false-success results literally impossible

All validation output MUST satisfy the following invariants:

1. Every gate reports concrete PASS / FAIL / SKIP with explicit reason text. A gate that just says "ok" is non-compliant.
2. Every SKIP reason MUST be mechanically distinguishable from a PASS. Summary counters track them separately.

3. FAILs are counted towards exit status. A script that hides FAILs in logs and returns 0 is a bug.
4. Meta-tests (§1.1) catch any gate that always PASSes regardless of condition.
5. Runtime results MUST be cross-referenced against host-side evidence (capture rig, screen recording, log analyzer, or equivalent) when the system's own introspection cannot directly verify the behaviour.

§ . NO BLUFF – positive-evidence-only validation

Every runtime test MUST satisfy ALL FIVE of the following non-negotiable constraints. A test that violates any one of them is a **bluff** and is forbidden from being committed to the test suite:

1. **Real ACTION**: the test MUST invoke at least one user-visible action via a project anti-bluff helper (e.g. `ab_send_action()` in shell-based suites). A test that records a PASS without ever calling such a helper is automatically failed by the summary function.
2. **State DELTA**: every PASS that claims to verify a feature MUST capture state BEFORE the action and AFTER, and assert the delta matches expectation.
`before == after` on success is a bluff.
3. **POSITIVE EVIDENCE**: the PASS condition MUST be POSITIVE evidence (a value match, a state delta, a captured frame analysis, a captured audio frame analysis), NEVER absence-of-error. If a probe returns nothing, that's a FAIL not a SKIP.
4. **UNIQUE EVIDENCE TOKEN**: tests that interact with mutable framework state MUST embed a per-run UUID into the action AND look for that token in the resulting state. Cached results predating the token cannot match — this defeats the stale-cache false-pass.
5. **Audio / video features REQUIRE captured evidence**: features that produce audio output MUST be validated by capturing the actual output via a capture pipeline (loopback, hardware capture rig, or equivalent) — NOT by metadata-only checks. Features that produce video output MUST be validated via frame capture + OCR or pixel-difference analysis.

Each consuming project SHOULD ship a shared anti-bluff helper library (typically `tests/lib/anti_bluff.sh` or equivalent) that codifies these rules; every new runtime test MUST source it and use its assertion helpers.

§ 7. Bleeding-edge ultra-perfection quality bar

The acceptance bar for shipping a project change is **not** "tests pass and source compiles" — that is the floor. The acceptance bar is:

The user, on a freshly-deployed system, sees the expected effect when they perform the documented action — confirmed end-to-end, with captured evidence, on the actual target environment.

The following are **forbidden**:

1. **"Source rebrand without shipped artifact rebuild"**: every source-side rename or rebranding MUST flow through to a shipped artifact that the runtime actually loads. Pre-build source-string gates are insufficient by themselves; shipped-artifact gates must accompany them.
 2. **"Tests pass but feature broken"**: this is the failure mode §7.1 exists to eradicate. Any test that PASSES despite the feature being non-functional is a critical defect; its mutation pair must FAIL on the breakage.
 3. **"Configuration-only tests"**: a test that only verifies a config file exists / has expected contents is downgraded to a pre-build gate. Runtime tests MUST exercise the actual user-visible behaviour and capture the result.
 4. **"It works on my workstation"**: every fix is validated on every supported target before any release. No exceptions.
 5. **"Will fix in next cycle"**: deferrals are documented in the changelog with concrete blockers + implementation plans, not just handwaved.
-

§ . Absolute codebase and data safety – zero risk, zero loss

Every destructive operation on the repository (history rewrite, force-push, branch deletion, bulk file removal, submodule de-init, pack repack that drops objects) is a **safety-critical** operation. Data loss from a wrong force-push is irreversible once the remote garbage-collects dangling objects. This section is non-negotiable.

§ . Mandatory safety protocol for destructive operations

Every destructive operation **MUST** execute, in strict order:

1. **Full backup before touching anything.** Hardlinked mirror of `.git` to a sibling backup directory:

```
cp -al .git <backup>/<repo>.git.mirror
```

This is near-instant on the same filesystem and uses zero additional disk. If `.git` is small, also create `git bundle create <backup>/main.bundle --all`.

2. **Record critical metadata into the backup dir:**

- `git show-ref > refs.txt`
- `git tag > tags.txt`
- `git submodule status > submodules.txt`
- `git rev-parse HEAD > head.txt`
- `git rev-parse HEAD^{tree} > head_tree.txt`
- `git ls-tree -r HEAD --full-tree | sha256sum > head_tree_content.sha256`
- `git ls-tree -r HEAD --full-tree > head_tree_listing.txt`

3. **Identify the target state:** the expected HEAD commit, tree hash, and tree content hash after the operation completes. The operation **MUST** preserve these unless it is explicitly changing HEAD content.
4. **Run the operation** (filter-repo, rebase, etc.). NEVER use `--no-verify` / `bypass-hooks` flags. NEVER auto-force on failure.

5. Post-operation verification gate (all MUST pass):

- HEAD tree hash matches pre-op target (content identical for non-content-changing rewrites)
- HEAD tree content sha256 matches pre-op sha256
- All tags from `tags.txt` are preserved and resolve to valid commits
- All submodule pointers from `submodules.txt` match current state
- Project's pre-build / pre-merge gates pass with zero FAIL and zero unclassified WARN

6. Only if every check in §9.1.5 passes may the operation be considered successful. Otherwise: restore from backup immediately (`rm -rf .git && mv <backup>/<repo>.git.mirror .git`).

7. ⁹ Then and ² only then may a force-push proceed — and still not without explicit user authorization per §9.2.

§ . Force-push requires explicit user authorization every time

Force-pushing (including `push --force` , `push --force-with-lease` , or any divergence-recovery code path inside an automated wrapper) overwrites the remote's view of history. On any shared branch this is irreversible.

Rules:

- Automated push wrappers MUST NOT auto-force-push as a fallback to any kind of failure (size rejection, divergence, server error, anything). The historical divergence-recovery code that silently force-resets the remote is forbidden and must be removed or gated behind an explicit flag.
- A human must affirmatively say "force-push" (or equivalent) in the session to authorize each force-push.
- Force-push to `main` (and any equivalent primary branch on owned submodules) is only authorized after the §9.1.5 gate has passed.
- Every force-push event is recorded in `docs/changeLogs/<tag>.md` under a "Force-push audit" section (target refs, backup location, validation gate output).

§ . Hardlinked backup is the standard – there is no excuse

Because hardlinked `.git` copies are near-instant and use zero additional disk (same-filesystem inodes), the cost of making a backup is trivial. Any destructive operation without a fresh hardlinked backup is a Constitution violation regardless of how small the repo is or how confident the operator.

§ . Commit-message audit trail for history rewrites

After any `git filter-repo` run (or equivalent), commit a record under `docs/changelogs/` (or `docs/history-rewrites/`) containing:

- What was stripped and why
- Pre-op `.git` size vs post-op `.git` size
- Pre-op HEAD commit vs post-op HEAD commit
- Pre-op HEAD tree hash vs post-op HEAD tree hash (MUST match for non-content-changing rewrites)
- Backup location

1 0

§ . Enforcement

Any commit, tag, or release that violates §1–§9 is non-compliant. The fix is to amend/revert and re-land with compliance. No exceptions for speed pressure. Data-safety violations (§9) and host-session-safety violations (§12) are treated as catastrophic and block the entire release cycle until fully remediated.

§ 11.4.107 . End-user quality covenant

§ 11.4.107 . End-user quality guarantee – forensic anchor (User mandate, – –)

The reason §7.1 + §8 exist — captured verbatim from the project owner so that future engineers and AI agents understand the historical failure mode the covenant exists to eradicate:

"We had been in position that all tests do execute with success and all Challenges as well, but in reality the most of the features does not work and can't be used! This MUST NOT be the case and execution of tests and Challenges MUST guarantee the quality, the completion and full usability by end users of the product!"

This is the canonical motivation for every gate, every mutation pair, every meta-test, every anti-bluff helper, every captured- evidence requirement, every multi-environment validation rule. The covenant exists because the project has historically shipped broken features behind PASS-ing tests, and that outcome is no longer tolerated.

Operative rule: the bar for shipping is not "tests pass" but "**users can use the feature.**" Every PASS in this codebase MUST carry positive evidence captured during execution that the feature works for the end user. Metadata-only PASS, configuration-only PASS, "absence-of-error" PASS, and grep-based PASS without runtime evidence are all critical defects regardless of how green the summary line looks.

Machine-created evidence at every gate (extension, 2026-08-15 — §11.4.262).

Every claim by the System that something WORKS, PASSES, IS VERIFIED, or IS PRODUCTION-READY MUST cite MACHINE-CREATED, MACHINE-VERIFIABLE evidence produced by the gate itself — not an operator's narrative, not "looks fine", not "no error was reported", not a prediction, not a guess, not an assumption.

§11.4.262 makes this obligation EXPLICIT at every gate seam and applies §11.4.107(10) self-validating analyzers UNIVERSALLY, not only to AV playback. The captured-evidence family (§11.4.5 / §11.4.69 / §11.4.107) is REFINED from

CAPTURED to CAPTURED-AND-MACHINE-VERIFIABLE. Prediction, guessing, and assumption about success are the exact anti-patterns this covenant forbids, at every layer of the System, from the first prompt of the session onward.

Propagation requirement: every `CLAUDE.md` and every `AGENTS.md` in the project tree (parent + every submodule recursively) MUST contain a clearly-titled "MANDATORY ANTI-BLUFF COVENANT — END-USER QUALITY GUARANTEE" section that quotes this user mandate verbatim and points back to this §11.4 as the canonical authority. The consuming project enforces this via a pre-build gate (typically `CM-COVENANT-PROPAGATION`).

§ 11.4 — FAIL-bluffs are equally forbidden

The covenant must be enforced symmetrically. The historical failure mode the covenant eradicates is the **PASS-bluff**: a green test result on a feature that doesn't work for users. But the inverse failure mode is just as toxic — the **FAIL-bluff**: a red test result caused by a script bug (shell crash, missing argument, regex error, malformed assertion) rather than by an actual product defect.

Examples of FAIL-bluffs:

- A test using `set -eu` calls `log_pass` without arguments → the helper references `$1` → script crashes with "1: parameter not set" → orchestrator records FAIL → engineer assumes product bug.
- A test pre-condition check uses an unbounded sed range → a comment line in the source matches the range terminator → the intended mutation never fires → audit reports PASS — bluff.

A FAIL produced by a test bug is no more useful than a PASS produced by a test bug — both mislead investigation, waste cycles, and ultimately let real defects ship undetected.

Operative rule (extended): every test MUST fail ONLY for genuine product defects. A test that crashes for a script-internal reason — undefined variable under `set -u`, unhandled error, malformed pipe, regex syntax error, division by zero, etc. — is a critical defect in the test suite itself and MUST be fixed at the source layer (helper library, shared lib, test source) so the failure mode disappears project-wide, not patched in individual call sites.

§ 11.4.3 – Recorded-evidence requirement

A test that emits PASS without **captured visual or audio evidence of the user-visible feature actually working on the screen the user would see** is a §11.4 PASS-bluff. Closing this gap requires a project-wide recording + analyzer infrastructure consisting of (at minimum):

- **Recording wrapper** that captures the user-visible output to time-synchronized media files. Multi-display systems record every output stream in parallel; sync metadata (host wall-clock at spawn) is written alongside the recording.
- **Action-timeline emitter**: every test emits structured timeline events (JSONL or equivalent) at every user-visible transition.
- **Frame / audio analyzer**: scans the recordings, extracts samples at a configurable interval, feeds them to recognizers (OCR for text overlays, speech-to-text for audio, pixel-diff for frame health), correlates against the timeline, and emits findings.

Closure criterion: every PASS for a user-visible feature **MUST** be cross-checked by the analyzer against the recording + action timeline. A PASS that lacks at least one matched timeline event in the analyzer findings is treated as a §11.4 PASS-bluff.

§ 11.4.4 – Per-environment-topology test dispatch

Tests that depend on environment topology (presence/absence of secondary displays, microphones, network interfaces, peripheral devices, cloud services, etc.) **MUST** detect topology at test entry and dispatch the topology-appropriate variant. A test that runs the **WRONG** variant for the actual topology and **PASSes** is a §11.4 PASS-bluff: it claims a feature works while never having exercised it on the configuration the user actually has.

A topology-touching test that does **NOT** have a dispatcher **AND** a per-topology variant is a §11.4 violation regardless of how it "passes" on any specific configuration. Single-test-multi-topology scripts **MUST** gate every topology-dependent assertion behind an explicit topology check that emits **SKIP-with-reason** if the required topology is absent — never **PASS-by-default**.

§ . . – Test-interrupt-on-discovery + retest-from-clean-baseline

A test cycle that *continues running past a freshly discovered defect* is itself a §11.4 PASS-bluff: it produces "all green" summaries that the operator might tag as a release while the codebase under test is known-broken at the moment those greens were recorded.

The mandatory protocol — non-negotiable:

1. **Interrupt immediately.** The moment a defect is re-discovered, re-produced, or newly identified during a test cycle, the cycle **MUST** stop. No "let it finish for the data" — the data is contaminated.
2. **Systematic debugging before any fix.** Identify root cause at source layer (not symptom), blast radius across related tests/features/subsystems, and regression-protection seam. Symptom patching is forbidden.
3. **Fix at root cause** per §11.4.1: source-layer fix, not symptom patch in the calling site. The fix **MUST** be safe (no new failure modes), minimal (no surrounding cleanup smuggled in), and non-regressive.
4. **Four-layer test coverage per fix.** Every fix lands with positive-evidence coverage in every applicable layer:
 - **Pre-build / pre-merge gate** (source-level catch)
 - **Post-build / packaging gate** (artifact-level catch)
 - **Runtime / on-device test** (end-user-behaviour catch)
 - **Meta-test paired mutation** (proves the gate is not itself a bluff)
5. **Documentation MUST be updated for every fix** — Issues → Fixed migration on closure, Applied Fixes table row, user-facing guides, architecture diagrams, per-version changelog.
6. **Rebuild the system.** Full build via the canonical build script. Skipping rebuild — even when the fix is host-only — invites stale-artifact contamination of the retest baseline.
7. **Re-deploy on every supported target.** A target left on the pre-fix build is a confounder.
8. **Repeat the full test suite** from the beginning on every target, sequentially per host memory cap. Partial reruns are a §11.4 violation.

¹¹ 9. **No-bluff certification per cycle.** Before tagging: meta-test harness returns all gates green AND every gate's paired mutation FAILs.

§ 11.4.5 — Captured-evidence quality analysis

§11.4.2 mandates *captured* evidence; §11.4.5 mandates the **content** of that evidence be analyzed for quality, not merely for presence.

Audio quality analysis — every audio test that **PASSes** **MUST** verify:

1. **Presence** — non-trivial RMS amplitude in captured output.
2. **Channel count** — analyzer reports the exact channel count the test claims (2.0 stereo, 5.1 surround, etc.). Stereo downmix in a "5.1 PASS" claim is a PASS-bluff.
3. **Sample rate + bit depth** — match the codec / pipeline under test.
4. **Glitch census** — underruns, XRUNs, dropouts above tolerance **MUST** be explicitly classified (PASS within budget, WARN above, FAIL on hard limits). Silently ignoring counts is a PASS-bluff.
5. **Coexistence-artifact census** — radio / IPC / shared-resource contention above tolerance.

Video quality analysis — every video test that **PASSes** **MUST** verify:

1. **Presence** — captured recording has non-zero size AND decoded- frame total > 0. A 0-byte file is the canonical PASS-bluff.
2. **Routing target** — analyzer confirms video appeared on the intended display / surface.
3. **Frame health** — drop count, jitter, freeze detection, tearing.
4. **Obstruction census** — OCR scan for hostile overlays (error dialogs, sign-in ¹¹ dialogs, geo-restriction overlays, paywall, etc).
5. **Resolution + codec** — captured dimensions match what the test claims.

§ 11.4.6 — No-guessing mandate

Tests, gates, status reports, closure narratives, commit messages, and any operator-facing text **MUST NOT** use words like likely , probably , maybe , might , possibly , presumably , seems , appears to , guess ,

seemingly , apparently , perhaps , supposedly , conjectured , or their synonyms when describing CAUSES of test failures, system behaviour, or fix effectiveness.

Either:

1. **Prove the cause** with captured forensic evidence (logs, kernel ring buffer, kernel ramoops, structured snapshots, OS journals, strace, etc.) and state it as fact, OR
2. **Explicitly mark UNCONFIRMED + PENDING_FORENSICS** with a tracked-task ID for follow-up.

Every "X likely caused Y" sentence in the codebase or documentation is a §11.4.6 violation^{11 4 7}. Every "appears to be benign" without a concrete forensic trace is a §11.4.6 violation.

§ . . – Demotion-evidence rule

A demotion from any FAIL classification (OPEN , POSSIBLE PRODUCT DEFECT , FAIL) to a lower-severity classification (INVESTIGATED , MITIGATED , RESOLVED , WORKING-AS-INTENDED) requires positive evidence captured under the SAME CONDITIONS that originally exposed the defect:

1. **Same target** — defect on Target-A needs Target-A evidence; cross-target claims need every target.
2. **Same build** — evidence captured BEFORE a relevant fix landed does not validate post-fix behaviour.
3. **Same cycle position** — a stress-soak FAIL is not refuted by an isolated-test PASS.
4. **Same load profile** — a contention-mode FAIL is not refuted by a quiescent-mode PASS.

"I cannot reproduce in isolation" is a HYPOTHESIS, not a finding. Per §11.4.6 it MUST be tagged UNCONFIRMED: until same-conditions retest produces positive evidence.

§ . . – Deep-web-research-before-implementation

Before designing a non-trivial fix, before implementing a new feature, before declaring an architectural choice — perform deep web research to verify the chosen approach is informed by current state-of-the-art. The research surface includes:

1. **Official documentation** of the platform, framework, language, or service.
2. **Vendor technical guides** for the hardware / cloud / library.
3. **Open-source codebases** that already solved analogous problems.
4. **Coding tutorials + technical articles**.
5. **Issue trackers** for the relevant projects.

Operative rule. A fix that re-invents a wheel (or reproduces a known-broken pattern) when the open-source community has already solved the problem is a §11.4 violation by omission.

Documentation requirement. Every non-trivial fix's commit message (or accompanying entry in the project's Issues / Fixed file) MUST cite the research sources that informed it: at least one external link OR the literal "NO external solution found — original work".

§ . . – Batch-source-fixes-before-rebuild

When working through a multi-defect closure cycle, all source-side fixes that DO NOT require runtime validation to design MUST be landed BEFORE the next artifact rebuild. The anti-pattern this mandate eliminates is `Fix A → rebuild → flash → cycle → fix B → rebuild → ...` which serializes operator time onto rebuild latency.

Exceptions where a fix DOES require interim rebuild MUST be documented in the fix's commit message as `REQUIRES_REBUILD: <reason>`. The default is the batch path; rebuild-now is the operator-authorized exception.

§ . . – Credentials-handling mandate

Forensic anchor — direct user mandate:

"Credentials or any secret and sensitive data MUST NOT leak! This MUST BE added as the mandatory constraint into all Submodules and main project's Constitution, CLAUDE.MD and AGENTS.MD."

All credentials, secrets, API tokens, passwords, phone numbers, OAuth refresh tokens, signing keys, encryption keys, and any other authentication-or-authorization-bearing data used by test scripts, build automation, or any project tooling MUST be handled per the following rules without exception:

1. **No tracked-file storage.** Credentials NEVER live in any file that git tracks. Placeholder templates with `<replace ...>` or `EXAMPLE_VALUE_DO_NOT_USE` placeholders ARE allowed in `.example` files committed to show operators what keys to populate.
2. **Repository-wide ignore.** `.gitignore` MUST exclude all credential-bearing paths:
 - `.env` , `.env.*` , `*.env` , `.<service>.env` , `*.<service>.env`
 - `scripts/testing/secrets/*` with `.example` + `README.md` exception
3. **Runtime-load-only.** Tests load credentials AT EXECUTION TIME from operator-populated files under `scripts/testing/secrets/` (or per-project equivalent). If the file is missing, the test SKIPS with "credentials not configured for " — it does NOT proceed without credentials and does NOT use defaults.
4. **No echo / no log.** Test scripts MUST NEVER print, log, or include credentials in error paths, stack traces, screen recordings, or output artifacts.
5. **Per-service file separation.** One file per service keeps blast radius small.
6. **Filesystem permissions.** `.env` files are `chmod 600` , parent directory is `chmod 700` .
7. **Rotation on suspected leak.** Rotate at provider, update local files, audit captured artifacts.

The consuming project's `.gitignore` MUST be verified before every commit:
`git ls-files --cached | grep -E "\\..env$" MUST show no real .env files (only .env.example ).`

§ . . .A – Pre-store credential leak audit (User mandate, – –)

Forensic anchor — verbatim user mandate (2026-05-17):

"Us these for all future testing (full automation testing) and make sure they are not leaking anywhere or get git versioned!"

[Discovered during execution: the operator-provided test credentials ALREADY existed in 5 tracked files dating to prior project cycles. §11.4.10 + §11.4.30 catch new commits but did not detect the pre-existing leak when the operator re-provided the same values for a new use.]

When an operator provides credentials, API tokens, signing keys, or any other secret material to be stored in the project's gitignored configuration (.env, scripts/testing/secrets/, secrets/, etc.), the agent or human storing them **MUST FIRST** execute a repo-wide audit for prior leaks of THOSE specific values **BEFORE** storing.

The audit **MUST**:

1. **Grep every tracked file** for the literal credential value(s). `git ls-files | xargs grep -l <value>` is the canonical form. Search **MUST** cover all file types — .md , .json , .yaml , .kt , .go , .py , .sh , .sql , build configs, docs.
2. **Grep the entire git history** for the literal credential value(s). `git log -S<value> --all --source --remotes` reveals historical commits even if the current tree is clean. A history-only leak is a §11.4.10 violation of equal severity to a tree-leak — anyone who pulled an older commit still has the value in their working copy of any branch they cloned.
3. **Report findings to the operator BEFORE storing the new value in any operator-controlled location.** The operator may then decide whether to (a) rotate the credential at the upstream provider before reusing, (b) accept the historical leak as a known-compromise and proceed, or (c) abort the storage. Storing the value without surfacing the audit is itself a §11.4 PASS- bluff at the security layer.

4. **If a leak is found:** open a forensic incident record per the project's §6/§7 sixth-law-incidents discipline (or equivalent), redact the literal values from all tracked files in-place (replace with `<redacted-per-§11.4.10>` placeholders), and record an OPERATOR ACTION REQUIRED for rotation per §11.4.10 sub-clause 7. The redaction commit lives on master IMMEDIATELY; the history-purge follow-up (git-filter-repo / BFG / equivalent + force-push to every upstream) requires explicit per-operation operator authorization per §9.2.
5. **Strengthen pre-push hook detection** when a previously-undetected class of literal credential is discovered. The pre-push hook's credential-pattern grep MUST be extended to catch the specific pattern class that escaped (literal usernames matched with literal passwords in the same file, well-known credential SHA-256 hashes, org-specific naming conventions for service accounts, etc.). The extension MUST land in the same commit as the redaction.

Pre-build gate (recommended) `CM-PRE-STORE-CREDENTIAL-AUDIT` (when implemented in consuming project) checks the commit message body of any commit that touches `.env` , `.env.example` , `scripts/testing/` `secrets/` , or equivalent for an `Audit-Pre-Stored:` stamp citing the audit's outcome. Paired mutation (§1.1) strips the stamp and asserts gate FAILs.

Composes with §11.4.10 (the core credentials mandate this extends), §11.4.30 (`.gitignore` + `no-versioned-artifacts` — overlapping enforcement surface), §11.4.6 (`no-guessing` — audit findings stated as fact or as `PENDING_FORENSICS`), §11.4.7 (demotion-evidence — "I think it's not leaked" without grep evidence is a guess). Classification: universal (per §11.4.17). No escape hatch.

§ . . – File-layout discipline

Project files MUST be organised by purpose, not by historical accident. Source code goes under canonical project roots. Tests go under canonical test directories. Logs and forensic artifacts go under operator-controlled directories (`~/Documents/` , `qa-results/` , `/tmp/` , OS-equivalent) — NEVER scattered at the repo root, NEVER inside working source trees, NEVER tracked unless they are reference assets explicitly required for evidence-replay.

Discovered drift is fixed by **moving** files to their canonical location and updating every caller — never by adding redirect shims or by leaving the misplaced copy "for backwards compatibility".

Carve-out (User mandate 2026-05-20). The 5 canonical tracker documents —

`docs/Issues.md` , `docs/Issues_Summary.md` , `docs/Fixed.md` , `docs/Fixed_Summary.md` , `docs/CONTINUATION.md` — sit at `docs/` root by design.

They are architectural constants of the project layout, analogous to AOSP's

`Makefile` , `Android.bp` , `OWNERS` files at repo root. Their location is encoded as literal path strings in §11.4.12 + §11.4.15 + §11.4.16 + §11.4.19 + §11.4.44 +

§11.4.53 propagation gates plus the helper-script constellation that regenerates

them (`generate_issues_summary.sh` , `generate_fixed_summary.sh` ,

`sync_issues_docs.sh` , `commit_docs.sh` , `colorize_progress_html.py` ,

`export_progress_docs.sh`). Moving them would require coordinated amendment

of those 6 sister anchors plus 5 pre-build gates plus ~20 helper scripts plus 42

consumer files (parent + 10 owned submodules + nested + HelixQA dependencies)

in a single PWU. Per §11.4.66, that scope is operator-blocked until explicitly

authorised. Audit-snapshot files (`docs/audit/anti_bluff_audit.md` , `docs/`

`audit/PRE_SONOS_TAG_READINESS.md` , `docs/audit/`

`D1_WIFI_FAIL_CLASSIFICATION.md` , plus any future audit snapshots) DO move

under `docs/audit/` per the §11.4.11 general principle — they are not part of the tracker constellation and carry no propagation-gate path-string burden.

§ 11.4.12 . 12 — Auto-generated docs sync mandate

Every auto-generated document (Issues_Summary, API reference, build manifest, etc.) MUST be regenerated in the **same commit** as any edit to its source. Stale auto-generated docs are §11.4 PASS-bluffs in document form: an operator reading them gets a divergent view of project state.

All three file types MUST stay in sync at all times — Markdown source + HTML export + PDF export (or equivalent multi-format output the project ships). Enforced by a pre-build gate that checks mtime ordering and content hash agreement.

§ . . – Out-of-band sink-side captured-evidence

Whenever a downstream consumer (HDMI sink, cloud service, monitoring target, downstream system) provides a network-accessible introspection API that reports what was actually received, the test suite **MUST** consume that report as captured-evidence for every test asserting end-to-end delivery.

The on-source-side view **ALONE** is insufficient — that is the exact "tests pass but the feature doesn't work" pattern §11.4 forbids. Sink-side reports **MUST** be:

- Identity-verified (MAC, serial, UUID match) before consumed as evidence.
- ¹Topology-dispatched (§11.4.3) — sink-probe-unreachable → SKIP, never FAIL.
- Cross-referenced with the on-source state at a matching wall-clock instant.

§ . . – Test playback cleanup mandate

A test that completes (PASS, FAIL, or SKIP) **MUST** leave the target in a quiescent state — no orphan playback, no orphan recording, no orphan capture, no orphan background process continuing after the test's last assertion. Tests are transient probes; their side-effects on shared target state **MUST NOT** outlive the test.

Mandatory protections:

1. Every test that issues a playback or capture command **MUST** issue the matching cleanup at test end (force-stop the app, stop the recorder, close the file descriptor, terminate the helper).
2. Cleanup is mandatory on **EVERY** exit path — PASS, FAIL, SKIP, error, trap, signal, parent-killed. Use shell `trap '<cleanup>' EXIT` or `try/finally`.
3. Tests **MUST** verify the cleanup succeeded via positive evidence (§11.4.5).
4. The orchestrator **MUST** run a post-test sanity check between tests and **FAIL** the just-completed test if it left orphan state. This blames the leaker, not the next-test victim.
5. No grace period for "the next test will clean it up" — that is precisely the §11.4 PASS-bluff pattern. Cleanup is the test's responsibility, not the next test's, not the orchestrator's.

§ 11.4.16 – Item-status tracking mandate

Every active item tracked in the project's Issues file MUST carry a `**Status:**` line within five lines of its heading. The status value is drawn from a closed set:

State	Meaning
Queued	In backlog, no work has started. Sub-state qualifier <code>Queued – BLOCKED on <reason></code> permitted.
In progress	Active investigation or coding underway.
Ready for testing	Source-side fix landed (pre-build / meta-test gates green), waiting for the next deploy cycle.
In testing	A deploy + cycle is in flight that exercises the fix; live captured-evidence is being collected per §11.4.2 + §11.4.5.
Reopened	Item failed its runtime test cycle and is back to active work. Carries a reference to the failure-cycle artefact.
Fixed (→ Fixed file)	Item closed: captured-evidence per §11.4.5 collected, runtime test PASSES, paired meta-test mutation FAILs, item migrated to the Fixed file.

Status MUST be updated as the item progresses. Every commit that changes an item's lifecycle state MUST update its Status line in the same commit.

All three ¹¹Issues / ⁴Issues_Summary / ¹⁶Fixed file types MUST be in sync at all times (Markdown + HTML + PDF).

§ 11.4.17 – Item-type tracking mandate

Every active item tracked in the project's Issues file MUST carry a `**Type:**` line within eight non-blank lines of its heading (the same window the §11.4.15 Status audit uses). The type value is drawn from a CLOSED set of three values:

Type	Meaning
Bug	Product defect / regression / user-visible broken behaviour. The product worked before (or was expected to) and now does NOT for the end user.
Feature	New capability not previously offered to end users — a new integration, new output, new probe, new bank, new architectural surface.
Task	Internal workstream — not user-visible. Refactor, infrastructure, documentation, audit, covenant clause, gate, mutation, propagation enforcement, host-session-safety hardening. The largest class by count; the lowest-stakes default when the classification is ambiguous.

The vocabulary is CLOSED — only Bug , Feature , Task . Any other value is a violation. When ambiguous, fall back to Task .

Type tells the operator WHAT KIND of work an item is — distinct from severity (HOW URGENT) and status (WHERE IN LIFECYCLE). The three axes together drive ownership routing, testing strategy, release-note framing, and changelog classification.

The Issues_Summary file MUST carry the Type column for every active item. All three file types (Markdown + HTML + PDF) MUST stay in sync under the same CM-DOCS-EXPORT-SYNC discipline as §11.4.12 + §11.4.15.

Pre-build gates CM-ITEM-TYPE-TRACKING (scans Issues file for **Type:** presence within the audit window) and CM-COVENANT-114-16-PROPAGATION (anchor presence in every project CLAUDE.md / AGENTS.md) enforce the mandate. Paired mutations in the project's meta-test prove neither gate is a bluff gate.

No escape hatch — items that genuinely don't fit MUST be re-classified as ^{11 4 17} INFORMATIONAL and excluded from the active-item set rather than carry a fabricated Type.

§ . . — Universal-vs-project classification of new rules (User mandate, - -)

Forensic anchor — direct user mandate (verbatim, 2026-05-14):

"Adding any new rules or mandatory constraints or anything relevant which should be added into Constitution, CLAUDE.MD, AGENTS.MD and relevant files MUST BE determined as reusable / universal VS project specific. If it is universal and reusable it MUST BE added into our root / main constitution Submodule, otherwise into our project level Constitution, AGENTS.MD and CLAUDE.MD (or Submodule if it is Submodule related)."

Before any new rule, mandatory constraint, covenant clause, gate declaration, or "MUST"-bearing statement is added to a project's Constitution / CLAUDE.md / AGENTS.md or to a child submodule's equivalents, the author MUST first classify it along this axis:

Classi- fica- tion	Definition	Destination
Univer- sal (re- usable)	The rule applies to ANY project — independent of language, framework, target platform, or domain. Examples: anti-bluff covenant, no-guessing mandate, credentials-handling, host-session safety, item-status tracking, multi-upstream push, file-layout discipline, deep-web-research-before-implementation.	This constitution submodule's <code>Constitution.md</code> + <code>CLAUDE.md</code> + <code>AGENTS.md</code> . Propagates to every consuming project via inheritance.
Pro- ject- specif- ic	The rule references a specific hardware target, vendor, model, SoC, vendor-fork, project-internal package name, deployment topology, or company-internal asset. Examples (illustrative — these are intentionally NOT in this universal file): the SoC's specific power-management quirks; a particular Android/iOS/Linux SDK behaviour; a specific app or service bundled with the project.	The project's own <code>Constitution.md</code> / <code>CLAUDE.md</code> / <code>AGENTS.md</code> (top-level), OR the affected submodule's equivalents (when the rule scopes to that submodule).

A rule that mentions hardware part numbers, vendor names, specific package identifiers, geographic regions, or company-internal asset names is **project-specific by definition** — it cannot be lifted into the universal layer without genericising those references first. A universal rule MAY reference patterns (e.g., "USB-HID multitouch panels") but MUST NOT name specific vendors.

Anti-bluff: the universal-vs-project classification is itself an auditable artefact. Every new rule's commit message MUST include a one-line classification statement (e.g.,

```
Classification: universal – applies to any project tracking items
through an Issues/Fixed lifecycle ). A commit that adds a rule without
classification justification is a §11.4.17 violation.
```

Authors who are uncertain SHOULD default to project-specific (narrower scope). Lifting a project-specific rule to universal later is cheap (one merge); generalising a universal rule that turned out to leak project-specific assumptions is expensive (every consumer must update).

Pre-build gate `CM-UNIVERSAL-VS-PROJECT-CLASSIFICATION` audits the last N^{11 4 20} commits for rule additions and asserts each one carries a classification statement. Paired mutation strips the classification literal^{2026 05 14} and asserts the gate FAILs.

§ . . – Subagent-driven-by-default mandate (User mandate, – –)

Forensic anchor — direct user mandate (verbatim, 2026-05-14):

"Make sure we ALWAYS WORK EVERYTHING subagents-driven if it is possible or applicable!"

When operating in a multi-agent runtime (Claude Code with subagents, Cursor with task-runners, Aider with sub-sessions, or any CLI tool that supports delegated execution), the primary agent MUST default to subagent delegation for any task that satisfies AT LEAST ONE of:

1. **Multi-step scope** — three or more discrete editing / file-creation / verification phases. Hand-off boundaries are where stalls happen; pre-planned subagent decomposition avoids them.

2. **Parallelisable work** — two or more independent tasks with no shared mutable state. Parallel subagents finish wall-clock-faster than sequential foreground work AND insulate the primary agent's context window from the volume of file reads.
3. **Long-running diagnostic loops** — repeated probe / re-test cycles, build-flash-cycle loops, soak tests. Subagents survive context compression boundaries that would otherwise truncate progress.
4. **Domain-specific or specialised workflows** — code review, security audit, infrastructure change review, performance triage, documentation propagation across N files. Specialised subagents (`code-reviewer` , `iac-reviewer` , `general-purpose`) bring pre-curated tool sets and disciplines the primary agent would re-derive from scratch.

The primary agent SHOULD only do foreground work when AT LEAST ONE of:

- The task is a single edit or single file read with no follow-up.
- The task requires conversational clarification with the operator mid-execution (subagents cannot ask the operator questions).
- The task is gating critical state (a commit, a push, a tag cascade) that must be sequenced before the next operator action.
- The task is so quick (under ~30 seconds of execution) that subagent spin-up overhead exceeds the work itself.

Anti-stall discipline. Subagents have watchdog timeouts in most runtimes (Claude Code: ~600 s of no-progress). When delegating, the parent agent MUST: (a) scope tightly (no "complete everything" prompts — break into 4-6 tightly-scoped tasks); (b) require **checkpoint commits** after each major task so partial progress is preserved even on stall; (c) explicitly prohibit destructive operations the subagent might attempt to "be helpful" (`git reset --hard` , `git push --force` , `--no-verify`); (d) provide explicit discipline pointers (§11.4.6 no-guessing, §11.4.10 credentials, §11.4.17 classification).

Anti-bluff applied. A subagent that returns "completed" without captured evidence in commits / outputs / artifacts is a §11.4 PASS-bluff at the multi-agent layer. The parent agent MUST verify subagent claims against repo state (`git log` , `git status` , post-completion gate runs) before treating the work as landed.

Coordination. When parallel subagents touch overlapping files, the parent agent MUST partition the work so subagents work on **non-overlapping files**. The parent's `commit_all.sh --auto-cascade` naturally bundles both subagents' dirty files into atomic commits via `git add -A` — exploit this for "forward-reference + provider" patterns where one subagent creates files another subagent references.

No escape hatch. Operating exclusively in the foreground when subagent delegation is feasible burns operator wall-clock time, risks context-window overflow on large tasks, and forfeits the parallelism that multi-agent runtimes were designed to provide. Pre-build gate `CM-SUBAGENT-DELEGATION-AUDIT` (when implemented per consuming project) scans recent session transcripts for multi-step foreground work that should have been delegated; paired mutation enforces the gate is not a bluff.

\$. . . — Script documentation mandate
(User mandate, — —)

Forensic anchor — direct user mandate (verbatim, 2026-05-14):

"Make sure that every single script inside the scripts dir (in all subdirs in depth) is fully documented and covered with full user manuals and user guides! ... Whenever is some bash script modified we MUST document it fully and create for it complete user guide(s) and manual(s) - if already exists make sure all of it is in sync and fully updated! No documentation ever can be out of sync with its codebase!"

Every Bash / shell / POSIX-sh script ANYWHERE in a project (anywhere under `scripts/` , `bin/` , `tests/` , library directories, `Upstreams/` , deployment hooks, CI helpers, etc. — depth-N recursive) MUST carry:

1. **In-source documentation block** at the top of the file:

- Purpose: one-sentence description of what the script does.
- Usage: complete invocation syntax with all flags, env vars, positional arguments, examples.
- Inputs: env vars, files read, command-line args, stdin.
- Outputs: files written, stdout/stderr behaviour, exit codes.

- Side-effects: anything that changes system state outside the script's own scratch directory.
- Dependencies: required commands (and how to install them) + required system state (e.g., "must be run as root", "must run from project root").
- Cross-references: companion guides under `docs/`.

2. **External user guide** under `docs/scripts/<script-name>.md` (Markdown).

Required sections:

- **Overview** — what the script does in plain English.
- **Prerequisites** — environment + dependencies.
- **Usage examples** — copy-paste recipes for the common cases.
- **Edge cases** — unusual invocations, error conditions, how to diagnose them.
- **Internal behaviour** — for advanced users / future maintainers: what the script does step-by-step.
- **Related scripts** — companion scripts and how they compose.
- **Last verified version / date** — when the doc was last reconciled against the script source.

Whenever the script source is modified, the in-source documentation block AND the external user guide MUST be updated in the SAME commit. A commit that touches a script but leaves its documentation unchanged (or worse, lets it drift) is a §11.4.18 violation.

Pre-build gate `CM-SCRIPT-DOCS-SYNC` walks every `*.sh` and `*.bash` under `scripts/` (or equivalent script directories), verifies a companion `docs/scripts/<name>.md` exists, AND verifies the doc was modified in the same commit as the script (by SHA cross-reference, or by `mtime ≥ script mtime` as a softer floor). Paired mutation strips the doc-sync invariant and asserts the gate FAILs.

No documentation ever can be out of sync with its codebase.

This mandate composes with §11.4.11 (file-layout discipline — docs live under `docs/`, scripts live under `scripts/`) and §11.4.12 (auto-generated docs sync — the on-disk Markdown + its rendered HTML/PDF must all stay synchronised).

§ . . — Fixed-document column-alignment mandate (User mandate, — —)

Forensic anchor — direct user mandate (verbatim, 2026-05-14):

"Make sure that Fixed document (all 3 file formats) is consistent and has the same structure and columns as Issues and Issues_Summary documents!
This has to be added as detail to Constitution, CLAUDE.MD and AGENTS.MD (constitution Submodule where these information shall exist)."

Every project that maintains an Issues / open-work tracker AND a Fixed / closed-archive tracker MUST keep the two structurally aligned along the same lifecycle axes — at minimum **Status** and **Type** — so an operator reading either gets a consistent view across the entire lifecycle (open → in progress → ready for testing → in testing → reopened → fixed).

Specifically:

1. **Per-entry annotation.** Every heading in the Fixed archive document MUST carry, within 8 non-blank lines of the heading, a `**Status:**` line and a `**Type:**` line — exactly as §11.4.15 + §11.4.16 require for the Issues tracker. For Fixed entries the Status value is drawn from the **closure-set**:
 - Fixed (→ Fixed.md) — fully verified on device with captured evidence per §11.4.5.
 - Fixed — pending device verification — source-side fix landed, regression-protection gate wired, awaiting the next firmware build + flash + cycle for captured evidence.
 - Fixed — RECLASSIFIED — item was reclassified during investigation (test-side defect, environment-only, etc.).

Type values follow §11.4.16: Bug | Feature | Task . Default on missing-annotation: Task .

2. **Auto-generated companion summary.** A Fixed_Summary.md MUST exist alongside Fixed.md , regenerated by an automation script (e.g. generate_fixed_summary.sh), and MUST mirror the column structure of Issues_Summary.md exactly:

	#	Level	Status	Type	One-line description
--	---	-------	--------	------	----------------------

The `Level` (severity) column uses the **original severity** the item had when it was open — so an operator can correlate a closed item with its original priority bucket. The same deterministic classification rules apply (REAL PRODUCT DEFECT / PARTIAL-or-WORKSTREAM-or-DEEP-DIVE / everything-else mapping to C / M / L).

3. **Three-format sync.** All three file types (`.md` , `.html` , `.pdf`) for BOTH the Fixed archive AND the Fixed_Summary MUST stay in sync at all times — same `.html` + `.pdf` export pipeline that §11.4.12 requires for Issues + Issues_Summary. The single-shot sync wrapper (e.g. `sync_issues_docs.sh`) MUST invoke the Fixed_Summary generator AND export all five doc variants (Issues, Issues_Summary, Fixed, Fixed_Summary, CONTINUATION) in one operation so they all carry the same mtime.
4. **Cross-lifecycle consistency.** A status line in Issues.md that reads `Fixed` (→ `Fixed.md`) is a **migration marker**, not a self-describing closure: when an Issues.md entry resolves, it MUST be moved to Fixed.md in the same commit, then disappear from Issues_Summary (because Issues_Summary only enumerates OPEN items) and appear in Fixed_Summary (because Fixed_Summary enumerates CLOSED items). Per §11.4.4 fix-closure protocol + §11.4.5 captured-evidence requirement, the migrated entry in Fixed.md MUST carry: closure cycle, closure commit SHA, captured evidence pointer, regression-protection gate name + paired mutation reference.

Pre-build gate `CM-FIXED-COLUMN-ALIGNMENT` (5+ invariants):

- `docs/Fixed_Summary.md` exists.
- `docs/Fixed_Summary.md` table header line contains both `Status` and `Type` columns (column-alignment with Issues_Summary).
- `mtime(Fixed_Summary.md) >= mtime(Fixed.md)` (regenerated after Fixed.md edits).
- Generator script exists (`generate_fixed_summary.sh` or equivalent integrated branch in `generate_issues_summary.sh`).
- Sync wrapper invokes the Fixed_Summary generator (grep for the script name in `sync_issues_docs.sh` or equivalent).

- HTML + PDF exports for both Fixed and Fixed_Summary exist (docs/Fixed.html , docs/Fixed.pdf , docs/Fixed_Summary.html , docs/Fixed_Summary.pdf).

Paired mutation: strip the Status column from Fixed_Summary.md table header (or rename the generator) → gate FAILs. Enforced by meta_test_false_positive_proof.sh .

Propagation. This anchor is a §11.4.17-classified **universal** rule — it composes naturally with §11.4.12 (auto-generated docs sync), §11.4.15 (Status tracking), §11.4.16 (Type tracking) and is mandatory across every consuming project's CLAUDE.md / AGENTS.md. The propagation gate CM-COVENANT-114-19-PROPAGATION (when implemented per consuming project) enforces the anchor's presence in every covenant file.

No escape hatch. A Fixed archive that lacks Status/Type columns or whose Summary companion drifts out of sync is the exact "different lifecycles report different facts" hazard §11.4 forbids — an operator reading Issues_Summary can no longer answer "where is item X now?" by cross-referencing Fixed_Summary because the schemas diverge. Non-compliance is a release blocker regardless of context.

2026 05 14

§ . . — Operator-blocked status + self-resolution exhaustion mandate (User mandate, - -)

Forensic anchor — direct user mandate (verbatim, 2026-05-14):

"Add into Issues and Issues_Summary new status: Operator blocked! We need another column with details in Issues document with exact details on how exactly is operator blocked issue! Before issue becomes operator blocked we MUST investigate in-depth all ways on how it can be resolved without operator's involvement - by you, or by activating any relevant mechanism for the resolution!"

Every project that maintains an Issues / open-work tracker **MUST** treat operator dependency as a **last-resort classification**, earned only after documented exhaustion of self-resolution paths. Routing an item to the operator without proving the agent + tooling cannot unblock it themselves is the §11.4 PASS-bluff pattern transplanted into the planning layer — the work stalls indefinitely while the agent claims "nothing more I can do."

1. Status vocabulary extension. §11.4.15's closed-set of Status values is extended with a seventh value, `operator-blocked`. The full closed-set becomes:

#	Status value	Meaning
1	<code>Queued</code>	Awaiting work to start.
2	<code>In progress</code>	Active agent / developer work.
3	<code>Ready for test-ing</code>	Source-side fix landed, awaiting verification.
4	<code>In testing</code>	Captured-evidence cycle running.
5	<code>Reopened</code>	Previously closed item resurfaced (regression).
6	<code>Operator-blocked</code>	NEW. Cannot proceed without an operator action that the agent + available tooling cannot perform.
7	<code>Fixed (→ Fixed.md)</code>	Closure migration marker per §11.4.15 + §11.4.19.

2. Self-resolution exhaustion mandate. BEFORE classifying any item as `operator-blocked`, the agent **MUST** verifiably exhaust every self-resolution path applicable to the project:

- **(a) CLI / ADB / SSH / API access** — can the work be done via any access the agent already has? (e.g. `adb shell pm clear`, `gh api`, `aws ssm`, `kubectl exec`, `psql`).
- **(b) Subagent delegation** — can a specialised subagent (per §11.4.20) reach further than the primary agent? (e.g. `general-purpose` for multi-step probing, `code-reviewer` for architectural blockers, `iac-reviewer` for infra access).

- **(c) Existing tooling** — is there a script / helper / library already in the repo that handles this case? (e.g. `scripts/testing/<helper>.sh` , project-specific automation).
- **(d) Captured fallback** — can the blocker be sidestepped with a synthetic event, test asset substitution, mock, or topology-aware SKIP per §11.4.3? (e.g. simulate the input, substitute a stand-in fixture, downgrade to a pre-build gate).
- **(e) Documentation + research** — per §11.4.8, has the agent consulted external sources (official docs, vendor guides, open-source codebases, issue trackers) for a self-resolution pattern the community has already solved?

Only AFTER each applicable path is **verifiably attempted and documented as exhausted** is `Operator-blocked` the correct classification. "I didn't try X because I assumed it wouldn't work" is a §11.4.6 no-guessing violation AND a §11.4.21 self-resolution violation.

3. Operator-Block-Details mandate. Every `Status: Operator-blocked` item MUST carry an `**Operator-Block-Details:**` line within 8 non-blank lines of its heading (mirroring the §11.4.15 Status placement pattern). The content MUST state ALL of:

- **WHAT** — the specific concrete action the operator must perform (verb + object + target system). Generic "the operator must investigate" is forbidden — name the action.
- **WHY** — the alternatives that were exhausted, enumerated. Each exhausted path from §11.4.21.2 above gets a one-line statement of what was attempted and why it could not unblock the item. "Subagent delegation: tried general-purpose with X scope — blocked because Y" is the bar.
- **UNBLOCK CONDITION** — the observable signal that the operator has completed the action (e.g. "operator confirms hardware rev-B installed", "operator pastes new API key into `scripts/testing/secrets/.foo.env` ", "operator force-pushes PR #123 merge to upstream"). Without this, the agent cannot automatically detect when to re-evaluate.
- **WHO** — the handle / contact / pointer to the document where operator-side details live if questions arise (e.g. "operator: @username", "see `docs/operator/<topic>.md` ", "vendor: support-id 12345"). Operator-blocked without WHO is operationally untrackable.

4. **Issues_Summary inclusion as sortable axis.** The auto-generated

`Issues_Summary.md` (per §11.4.12 + §11.4.15) MUST include `Operator-blocked` as a first-class Status value in the table — the column header is unchanged but operators can sort / filter on it deterministically. The summary generator MUST also emit a count of `Operator-blocked` items in any rollup line / header.

5. Periodic re-evaluation requirement. Items in `Operator-blocked` status MUST be re-evaluated every Nth tag-cycle (project-defined, recommended every 3rd tag cycle). Operator dependencies change over time: CI may have been added, hardware may have been swapped, automation may have matured, vendor APIs may have unlocked. An item that was correctly Operator-blocked at tag T may be self-resolvable at tag T+3. Re-evaluation MUST follow the same §11.4.21.2 exhaustion checklist and document each re-attempt.

6. Anti-bluff layer. A fake `Operator-blocked` (classification applied without exhausting §11.4.21.2 alternatives) is a §11.4 covenant violation at the planning layer — equivalent severity to a PASS-bluff at the testing layer. The agent's commit message introducing or maintaining an `Operator-blocked` classification MUST contain evidence of self-resolution attempts (the explicit "Attempted: a — exhausted because X; b — exhausted because Y; c — exhausted because Z" form).

7. **Pre-build gates (recommended, per consuming project):**

- **CM-ITEM-OPERATOR-BLOCKED-DETAILS** — every heading whose Status line equals `Operator-blocked` has an `**Operator-Block-Details:**` line within 8 non-blank lines. Paired mutation strips the details line → gate FAILs.
- **CM-OPERATOR-BLOCKED-SELF-RESOLUTION-AUDIT** — every NEW `Operator-blocked` item introduced in a commit carries an "Attempted: ..." audit trail (either in the Issues.md entry body OR in the commit message). Paired mutation introduces an `Operator-blocked` item without the audit trail → gate FAILs.

Propagation. This anchor is a §11.4.17-classified **universal** rule — it composes with §11.4.15 (Status tracking), §11.4.16 (Type tracking), §11.4.19 (Fixed-document column alignment), §11.4.12 (auto-generated docs sync), §11.4.20 (subagent delegation as a self-resolution path), §11.4.8 (research-before-implementation as a self-resolution path), §11.4.6 (no guessing about what the operator "would have to

do"). Propagation gate `CM-COVENANT-114-21-PROPAGATION` (when implemented per consuming project) enforces the anchor's presence in every `CLAUDE.md` / `AGENTS.md` across the parent + every owned submodule + every dependency.

No escape hatch. Items that legitimately need operator intervention are real and unavoidable — hardware swaps, vendor contracts, physical-world inputs, account credentials, irreversible business decisions. But the classification must be **earned** via documented exhaustion. An `Operator-blocked` heading without an audit trail of self-resolution attempts is a §11.4 release blocker regardless of context.

2026 05 14

§ . . — Document-sync commit discipline (User mandate, — —)

Forensic anchor — direct user mandate (verbatim, 2026-05-14):

"For Issues, Issues_Summary and Fixed docs (and all its exports) we MUST commit and push (only them) as soon as they are updated (synced). Continuation doc and other similar / relevant documentation MUST be committed and pushed too! Like this we will always have up to date status of working items without need to do full commit and push of everything if that is not possible to do!"

Classification: §11.4.17-classified **universal** — every project that tracks work items through an Issues / Fixed lifecycle has the same problem, so the discipline lives in the Constitution and every consuming project implements its own wrapper.

The defect this anchor closes. Until this anchor existed, the only way to push a status update was the full-tree commit wrapper. When the working tree had unrelated heavy churn (large submodule diffs, in-flight rebase, partial-network conditions, ongoing build), operators had two bad options: (a) wait until the heavy work finishes — which can be hours — leaving the doc-status stale to every other agent; (b) skip the doc-status update entirely — which is a §11.4.4 documentation-drift violation. Neither option is acceptable.

The mandate. Every project under this Constitution that ships a status-tracking doc set MUST provide a **lightweight commit path** distinct from the full-repo commit wrapper. The lightweight path stages, commits, and pushes **only** the status-tracking doc set:

1. The active issue tracker (`docs/Issues.md` in a consuming project's layout; the project's equivalent file elsewhere) and its HTML + PDF exports.
2. The auto-generated issues summary (`docs/Issues_Summary.md`) and its HTML + PDF exports.
3. The fixed-bug archive (`docs/Fixed.md`) and its HTML + PDF exports.
4. The auto-generated fixed summary (`docs/Fixed_Summary.md`) and its HTML + PDF exports.
5. The cross-session continuation document (`docs/CONTINUATION.md` , per §12.10) and its HTML + PDF exports.
6. Any auto-generated audit artifact (`docs/audit/anti_bluff_audit.md` or equivalent) that pairs with the above.

The lightweight wrapper MUST:

1. **Auto-invoke the project's export-regeneration pipeline first** — so Markdown + HTML + PDF are all synchronised before push. (Skipped only when an explicit `--no-sync` flag is passed AND the operator just ran the pipeline manually.)
2. **Stage ONLY the doc-set files** — `git add` of an explicit file list, NEVER `git add -A` (which would catch unrelated WIP churn).
3. **Use a separate flock** disjoint from the full-tree commit wrapper's lock — so the two can run in parallel without contention on disjoint file sets.
4. **Push to every parent-repo remote** (reusing the project's push-all driver where available).
5. **Exit code semantics** — `0` success, `1` validation failure, `2` lock contention, `3` nothing-to-commit (informational, not an error — so the wrapper can be wired into automation that tolerates no-op outcomes).

The lightweight wrapper MUST be invocable either standalone OR via a flag on the full-tree wrapper (e.g. `commit_all.sh --docs-only`) so operators have a single mental model.

§9 preflight inheritance. The lightweight wrapper inherits the parent project's §9 preflight discipline — if `meta_test_*` is mid- mutation or `*.mut.bak` files exist, the wrapper refuses to run, same as the full-tree wrapper.

Pre-build gates (recommended, per consuming project):

- **CM-COMMIT-DOCS-EXISTS** — verifies the lightweight wrapper exists + is executable, its external user guide (per §11.4.18) exists, the full-tree wrapper advertises the delegation flag, and the wrapper's doc-set enumeration lists $\geq N$ entries (N depends on project — a consuming project's set is 15). Paired mutation strips the doc-set array → gate FAILs.

Propagation. Composes with §11.4.12 (export-sync invariant — HTML + PDF in sync with Markdown after every edit), §11.4.15 (item status tracking — Status lines must be visible quickly), §11.4.18 (every script ships with an in-source doc block + an external user guide), §12.10 (CONTINUATION document maintenance — the lightweight wrapper is one of the mechanisms that keeps CONTINUATION current).

No escape hatch. The discipline exists because doc-status drift is itself a §11.4 PASS-bluff pattern at the documentation layer: a green-looking `Issues_Summary.html` that doesn't reflect Markdown reality is the same class of defect as a passing test that doesn't reflect runtime reality. Operators who can't run the lightweight wrapper for some reason **MUST** run the full-tree wrapper and accept the heavier cost — the doc-status drift is non-negotiable.

2026 05 14

§ . . — Visual-cue & grouping mandate
for Issues docs (User mandate,
- -)

Forensic anchor — direct user mandate (verbatim, 2026-05-14):

"We MUST make sure that Issues, Issues_Summary and Fixed docs and its exported files (PDFs and HTMLs) have one small improvement: we MUST introduce some background coloring for cells of item type and status! Also, we MUST group items by the status! ... Base colors for types would be: bug - background pale red, task - background pale blue, feature - background pale yellow. However, changing the status affects the color of the cell of the type and the status cell background color: queued - no color, no effect, fixed - both cells pale green, in progress - only status cell is pale green, reopened - status cell is pale red, blocker - status cell is red (not pale, however MUST BE still readable), anything else please follow the same logic!"

Classification: §11.4.17-classified **mixed** — the discipline (visual cues + status grouping for tracked-item docs) is universal across every project that maintains an Issues / Fixed lifecycle. The implementation (CSS classes, HTML post-processor, weasyprint integration) is project-specific because each project's export toolchain differs (pandoc, asciidoctor, sphinx, etc.).

The defect this anchor closes. A long, multi-column tracked-item table where every row looks identical at first glance is a §11.4 operational-readability defect — operators have to *read* every Status column to identify what is queued vs in-progress vs fixed. With dozens to hundreds of items, this is slow enough that operators inevitably skim, and the doc effectively becomes write-only. Color cues + status grouping make at-a-glance triage possible.

The mandate. Every project under this Constitution that ships tracked-item docs (per §11.4.15) MUST apply visual-cue coloring AND status grouping to those docs' HTML + PDF exports. The Markdown source stays free of color noise (Markdown is the canonical text), but the HTML + PDF exports MUST be enriched in a post-processing stage:

Type cell base colors:

Type	Background
Bug	pale red (#FFE5E5 or equivalent — ~5% red saturation)
Task	pale blue (#E5F0FF)
Feature	pale yellow (#FFF8DC)

Status cell colors (override the Status column; Fixed also overrides Type cell):

Status	Status cell	Effect on Type cell
Queued	no color (transparent)	keeps base type color
In progress	pale green (#D4F1D4)	keeps base type color
Ready for testing	pale green-blue (#C6EAEF)	keeps base type color
In testing	pale green-yellow (#E7F4D6)	keeps base type color
Reopened	pale red (#FFCCCC)	keeps base type color
Fixed	pale green (#A8E6A8)	also pale green
Operator-blocked / Blocker	vibrant red (#FF4444) with white text for readability	keeps base type color

Status grouping: items in the rendered table **MUST** be grouped by Status, emitting an H3-class heading for each Status group with the item count. Group order (most-actionable first; closed last):

1. Operator-blocked / Blocker
2. In testing
3. Ready for testing

4. In progress
5. Reopened
6. Queued
7. Fixed

This ordering surfaces the items requiring operator attention at the top of the doc; closed/archived items sink to the bottom.

Print-fidelity requirement. The CSS MUST declare `print-color-adjust: exact` (or the project's renderer's equivalent) so PDF exports preserve the cell backgrounds. A weasyprint default-rendered PDF that strips colors is a §11.4 export-drift defect — HTML and PDF MUST be visually equivalent.

Pre-build gates (recommended, per consuming project):

- **CM-DOC-COLOR-GROUPING-DISCIPLINE** — verifies the project's CSS carries the type+status class set, the post-processor exists and is executable, the sync wrapper invokes it, AND after sync the rendered HTML carries the grouping-heading + per-cell color classes (the captured-evidence proof — without this last invariant a stub post-processor could silently pass the static gate). Paired mutation strips a Status class literal from CSS → gate FAILs.

Propagation. Composes with §11.4.12 (export-sync invariant — HTML + PDF in sync with Markdown after every edit), §11.4.15 (item status tracking — every active item carries a Status line from the closed-set vocabulary), §11.4.19 (Fixed_Summary column-aligned companion), §11.4.22 (lightweight commit path for doc-set updates).

No escape hatch. Color cues are not optional polish — they are the operational-readability seam that converts a long tracked-item table from a write-only Markdown blob into a self-triaging document.

2026 05 14

§ . . — Build-resource stats tracking
mandate (User mandate, — —)

Forensic anchor — direct user mandate (verbatim, 2026-05-14):

"We need to incorporate through the Containers Submodule we use to run our System build Container the following mechanism safely: during its working lifetime we MUST track in proper Markdown document exported into PDF and HTML stats on System resources use! What was the peak and the minimum EVER the Container / System building process used! ... In top of the document we MUST have this ever values. ... After this we MUST present properly sorted divided by version tags names all build processes we were running with notes if the process completed with success or not, and resources usage during the process! Besides min and max values ever and per build iteration we MUST have for each resource we have average values — average memory usage, average CPU usage, and all other parameters IO and other resources!"

Classification: §11.4.17-classified **mixed** — the discipline of tracking host-side build-resource usage with min / max / mean / p95 per-build PLUS ever-values across all builds in a versioned Markdown + HTML + PDF triple is universal across every long-build project (AOSP, kernel, large monorepo, ML model training, multi-hour data pipelines). The implementation (registry path, monitor script name, exporter wiring) is project-specific.

The defect this anchor closes. Long builds (multi-hour AOSP builds, ML training runs, large monorepo CI) consume highly variable resource profiles over their lifetime — soong_build's 33 GB peak coexists with quiescent linker-only stretches at <2 GB. Without per-build sampling + ever-values, operators have NO empirical basis to (a) size containers correctly, (b) detect resource regressions across versions, (c) prove which build iteration was the OOM-killer's trigger, (d) reason about parallelism caps (cf. §12.7 -j2 hard cap whose forensic basis would have been one full Stats.md generation earlier had this discipline been in place). Memory pressure debugging without time-series data is the bluff this anchor forbids.

The mandate. Every project under this Constitution with a build exceeding **1 minute wall-clock** MUST:

1. Run a host-side resource sampler for every build. The sampler runs as a sibling background process (not inside the build's own process tree), reads `/proc/meminfo` + `/proc/loadavg` + `/proc/stat`

- `/proc/diskstats` (or platform equivalents on non-Linux) at a fixed interval (recommended 5 s default — fine-grained enough to catch seconds-long peaks; coarse enough to keep itself <5% CPU and <50 MB RSS). Samples written as TSV.

2. Compute per-build summaries on stop. Per metric (memory used, CPU%, load average, disk read/write): **min, max, mean, p95**. p95 is the 95th-percentile sample value — required because mean alone hides extended-peak behaviour (a 2 h build with one 33 GB soong_build peak has mean $\approx$ 18 GB but p95 $\approx$ 32 GB; the latter is what operators must size for).

3. Append per-build summaries to a registry. One TSV row per build (build_id, status SUCCESS/FAIL/UNKNOWN, start/end timestamps, per-metric min/max/mean/p95, total disk read/write). The registry is the single source of truth — the Markdown report is derived.

4. Maintain ever-values across ALL builds. Min / max / mean across every tracked build, computed on every Markdown regeneration from the registry. Surface at the **top** of the report per User mandate verbatim.

5. Markdown + HTML + PDF triple. Stats.md (auto-generated), exported to Stats.html + Stats.pdf through the project's normal export pipeline. Triple stays in sync per §11.4.12. Committed via the project's lightweight doc-sync wrapper per §11.4.22.

6. Sort per-build entries by recency (most recent first) AND, where applicable, group by version tag. Operators reading the report want the latest build's profile at the top and the historical sweep underneath.

7. Track status per build. SUCCESS / FAIL / UNKNOWN — and if FAIL, capture the reason (exit code, OOM, ANR, etc.). A FAIL build's resource profile is the most operationally interesting one in the report — never silently drop FAIL entries.

8. Performance budget for the sampler itself. <50 MB RSS, <5% CPU. The monitor MUST NOT itself contribute to the resource pressure it is measuring (Heisenberg-class observer effect at scale). Pure /proc reads + awk arithmetic is the reference implementation; vmstat / pidstat acceptable; psrecord / heavy-Python-import samplers are forbidden.

9. Survive build failures. The stop hook MUST be called from both the success AND failure paths of the build wrapper so FAIL builds still produce a row. Hooking only the success path is a PASS-bluff at the telemetry layer.

Pre-build gate (recommended, per consuming project):

- **CM-BUILD-RESOURCE-STATS-TRACKER** — verifies (a) the monitor script exists + executable, (b) the Stats.md target file exists, (c) the build wrapper invokes start AND stop, (d) the doc-sync wrapper enumerates the Stats.{md,html,pdf} triple, (e) the exporter advertises the build-stats slug. Paired mutation hides the monitor script aside → gate FAILs.

Propagation. Composes with §11.4.12 (export-sync invariant — HTML + PDF in sync with Markdown), §11.4.18 (script-documentation discipline — the monitor ships with an external user guide), §11.4.22 (lightweight doc-sync commit wrapper enumerates the Stats.{md,html,pdf} triple), §12.6 / §12.7 / §12.9 (host-safety + containerized-build envelope whose forensic anchors are the empirical motivation for this telemetry).

No escape hatch. Build-resource tracking is not optional polish — it is the operational-evidence seam that converts post-mortem "the build OOM'd, no idea why" into "build X at git-SHA Y peaked at Z GB at minute N, which is W% above the rolling p95". The discipline pays for itself the first time it diagnoses a build-resource regression.

§ . . — Full-Automation-Coverage
Mandate (User mandate, — —)

Forensic anchor — verbatim user mandate (2026-05-15):

"Make sure that every feature, every functionality, every flow, every use case, every edge case, every service or application, on every platform we support is covered with full automation tests which will confirm anti-bluff policy and provide the proof of fully working capabilities, working implementation as expected, no issues, no bugs, fully documented, tests covered! Nothing less than this does not give us a chance to deliver stable product! This is mandatory constraint which **MUST BE** respected without ignoring, skipping, slacking or forgetting it!"

Operative rule. For every consuming project under this Constitution, no feature, functionality, flow, use case, edge case, service, or application on any supported platform may be considered **deliverable** until it is covered by automation tests that collectively prove six invariants:

1. **Anti-bluff posture (per §7.1 + §11.4):** every assertion carries captured runtime evidence; metadata-only / configuration-only / grep-based / absence-of-error PASSes are forbidden.
2. **Proof of working capability:** the user-visible behaviour is exercised end-to-end on the target platform topology (per §11.4.3), not in a mock or in-memory facsimile.
3. **Working implementation as expected:** assertions match the product's documented promise (in user manual, README, specs), not an implementation-detail back-door.
4. **No issues, no bugs:** the test suite is the canonical seam for surfacing defects — passing without producing defect signals means defects do not exist or have been previously surfaced, tracked (per §11.4.15 / §11.4.16), and closed.
5. **Fully documented:** the feature has a user-facing doc entry (per §11.4.18 for scripts, and project-level user-manual coverage for everything else); the doc is kept in sync with the tests (per §11.4.12).
6. **Tests-covered (the four-layer floor per §1):** pre-build presence-of-change gate, post-build artifact-shipped gate, runtime / integration / on-device gate, AND meta-test paired mutation that proves the runtime gate catches the break.

Cross-cutting reach. This mandate is **universal** — it applies to every project consuming this Constitution, every feature surface they expose (HTTP endpoints, CLI commands, slash commands, IPC channels, plugins, hooks, MCP servers, agents, providers, integrations, GUI flows, mobile flows, scheduled jobs), and every supported platform (Linux, macOS, Windows, iOS, Android, AuroraOS, HarmonyOS, embedded, containers, headless servers, kiosk, etc.). A project that ships a feature without satisfying all six invariants is **not delivering a stable product**, irrespective of how green its summary line looks.

Coverage audit. Consuming projects **MUST** publish a coverage ledger (matrix of: feature × platform × invariant-1..6 × status) that is regenerated as part of the release-gate sweep. The ledger itself is documented (per §11.4.18 / §11.4.12) and committed via the §11.4.22 lightweight doc-sync wrapper. Gaps in the ledger (UNCONFIRMED: / PENDING: / BLOCKED: cells) **MUST** cite a tracked work item per §11.4.15 + §11.4.16; rows that quietly omit a platform are §11.4.25 violations.

Composition. §11.4.25 explicitly stacks on top of §1 (four-layer test-coverage floor), §1.1 (false-positive immunity), §7.1 (positive-evidence-only validation), §11.4.1 (FAIL-bluffs forbidden), §11.4.2 (recorded-evidence requirement), §11.4.3 (per-environment-topology dispatch), §11.4.6 (no-guessing), §11.4.15 / §11.4.16 (status + type tracking), §11.4.17 (universal-vs-project classification — this rule is universal), §11.4.18 (script documentation), §11.4.20 (subagent delegation when the work is multi-step), §11.4.22 (lightweight doc-sync). It does NOT supersede them; it forecloses the loophole "the feature works locally for me, ship it".

Classification: universal (per §11.4.17). No escape hatch. Severity-equivalent to a §11.4 PASS-bluff at the release-gate layer — a project claiming "Done" without honoring §11.4.25 is making a false claim regardless of how the work-item tracker labels the row.

2026 05 15

§ . . – Constitution-Submodule Update
Workflow Mandate (User mandate,
– –)

Forensic anchor — verbatim user mandate (2026-05-15):

"Every time we add something into our root (constitution Submodule) Constitution, CLAUDE.MD and AGENTS.MD we MUST FIRST fetch and pull all new changes / work from constitution Submodule first! All changes we apply MUST BE committed and pushed to all constitution Submodule upstreams! In case of conflict, IT MUST BE carefully resolved! Nothing can be broken, made faulty, corrupted or unusable! After merging full validation and verification MUST BE done!"

Operative rule. Before ANY agent or operator modifies the `constitution/Constitution.md` , `constitution/CLAUDE.md` , or `constitution/AGENTS.md` files of a project that consumes this Constitution as a submodule, the agent or operator MUST execute the following pipeline in order, with NO step skipped:

1. **Fetch + pull first.** From inside the `constitution/` submodule worktree, run `git fetch <every-configured-remote>` followed by `git pull --ff-only origin <branch>` (or `--rebase` if non-FF-mergeable; **never** `--strategy=ours` / `--allow-unrelated-histories` without explicit operator authorization). The submodule MUST be at upstream tip BEFORE any local edit is applied.
2. **Apply the change.** Edit the relevant file(s). The edit MUST classify itself per §11.4.17 (universal vs project-specific) — only universal additions belong in the constitution submodule; project-specific clauses belong in the consuming project's own governance files. The edit MUST include the verbatim user mandate (if it originated from one) as a forensic anchor.
3. **Validate before commit.** Run the constitution submodule's `meta_test_inheritance.sh` (or equivalent) to confirm the inheritance chain still resolves. Verify no governance file was left with merge-conflict markers (`<<<<<<<` , `=====` , `>>>>>>>`). Verify all three governance files (Constitution + CLAUDE + AGENTS) cross-reference the new clause consistently.
4. **Commit + push to ALL upstreams.** Stage only the governance files (NEVER `git add -A` inside the constitution submodule — stray local artefacts MUST NOT enter governance). Commit with a message that cites the user mandate (verbatim quote) + the classification line per §11.4.17. Push to **every** configured

remote of the constitution submodule. A commit that lives on one upstream but not others is a §11.4.26 violation equivalent to a §2.1 multi-upstream-push violation.

5. **Conflict resolution.** If `pull --ff-only` reports non-fast-forward, the merge MUST be performed carefully: inspect both sides, preserve the union of governance content (no clause silently dropped), re-classify per §11.4.17, validate per step 3. Force-push to "make conflicts go away" is FORBIDDEN (§9.2). Nothing about the constitution may be broken, made faulty, corrupted, or rendered unusable by the merge.
6. **Post-merge validation + verification.** After the push lands, re-clone (or `git submodule update --remote --init`) in a throwaway worktree and re-run the consumer project's inheritance-cascade verifier (e.g. `scripts/verify-governance-cascade.sh`) to confirm the new clause reaches every owned submodule per CONST-047. Any cascade gap MUST be closed in the same change-window.
7. **Update the consuming project's pointer.** The consuming project's `.gitmodules`-tracked submodule pointer MUST be bumped to the new constitution HEAD in the SAME commit as any downstream cascade work; out-of-sync submodule pointers are a §11.4.26 violation.

Operational scope. This workflow applies regardless of who initiates the change (operator, primary agent, subagent per §11.4.20, automated lint suggestion). The workflow CANNOT be shortcut by "I'll fetch later" or "I'll push to the other upstreams in the next commit". The constitution is the **single source of truth** for every project that imports it; allowing it to fragment across upstreams is the structural equivalent of a §11.4 PASS-bluff at the governance layer.

Cross-cutting reach. §11.4.26 composes with: §2 (single- entypoint commit wrapper), §2.1 (multi-upstream push), §3 (submodule changes propagate through submodule commits first), §9.1 / §9.2 (data-safety + force-push authorization), §11.4.17 (universal-vs-project classification), §11.4.22 (lightweight doc-sync), §11.4.25 (full-automation coverage — the post-merge validation step is itself a form of automation coverage). It does NOT supersede them.

Classification: universal (per §11.4.17). No escape hatch. A constitution-submodule change that violates §11.4.26 is a release blocker for every consuming project, equivalent in severity to a force-push without §9.2 authorization.

§ . . — Submodule-Dependency-Manifest Mandate (User mandate, — —)

Forensic anchor — verbatim user mandate (2026-05-15):

"We MUST HAVE mechanism for each Submodule to determine / know what are its Submodule dependencies so new projects or palces we are incorporate them can add these Submodules to the project root and make them available! Suggested idea is configuration file with expected Submodules Git ssh urls perhaps? New project can read it, and recursively add each Submodule to the root of the project and install / expose it to everyone. This MUST be analyzed and applied. We MUST apply the best strategy for this which can be easily executed just by following our root Constitution, AGENTS.MD and CLAUDE.MD! Process this, extend out Constitution, AGENTS.MD and CLAUDE.MD with mandatory instructions and then process project root and all Submodules deep recursively so proper configuration Submodules dependency files are created! Document EVERYTHING and cover with all supported test types! Any kind of bluff is strictyl forbidden! All wotk MUST come with mechanism for validation and verification by creating proper proofs for all critical points!"

Operative rule. Every owned-by-us submodule MUST ship a machine-readable, version-controlled **dependency manifest** at the canonical path `helix-deps.yaml` (or `helix-deps.json` / `helix-deps.toml` if the submodule's ecosystem strongly prefers a different serialisation — but only one canonical file per submodule, and its name MUST appear in the submodule's `README.md` + governance trio).

Manifest schema (CONST-051(C)-aligned):

```

# helix-deps.yaml
# Declares this submodule's own-org dependencies. Consuming projects
# read this manifest recursively and add each declared dependency to
# their root at <root>/<name>/ or <root>/submodules/<name>/ per
# CONST-051(C). Nested own-org submodule chains are FORBIDDEN — this
# manifest is the bridge. (Sole exception per §11.4.28(C): the
# constitution submodule MAY host depth-1 reusable-engine submodules
# under constitution/submodules/<name>/, each shipping its own
# helix-deps.yaml — never a recursive chain.)

schema_version: 1
deps:
  - name: Challenges                                # canonical name (will become snake_case per §11.4.29)
    ssh_url: git@github.com:vasic-digital/Challenges.git
    ref: main                                       # branch or pinned tag
    why: "Cross-cutting Challenges + Panoptic browser harness"
    layout: flat                                  # 'flat' = <root>/<name>/; 'grouped' = <root>/submodules/
                                              <name>/
  - name: HelixQA
    ssh_url: git@github.com:HelixDevelopment/HelixQA.git
    ref: main
    why: "Anti-bluff QA orchestration + autonomous-session driver"
    layout: flat

transitive_handling:
  # Each declared dep itself ships a helix-deps.yaml. The incorporator
  # tooling MUST recurse — top-level project gets the union of all
  # transitively-declared deps, flattened to root.
  recursive: true

  # When two submodules declare the same dependency at different
  # refs, CONFLICT — operator MUST resolve before incorporation
  # proceeds. The incorporator aborts on conflict (never silently
  # picks one).
  conflict_resolution: operator-required

language_specific_subtree: false    # set true for Android/Kotlin/Apple
                                   # roots (per §11.4.29 exception);
                                   # excludes inner subtree from
                                   # snake_case enforcement.

```

Tooling contract. A consuming project **MUST** be able to bootstrap its entire own-org dependency graph by:

1. Running `incorporate-submodule <ssh-url>` (canonical name; lives in the constitution submodule's `scripts/` or in a parent project's bin path).
2. The script:
 - Adds the supplied submodule at its declared canonical path (per CONST-051(C) `flat` vs `grouped` layout).
 - Reads the newly-added submodule's `helix-deps.yaml`.
 - For each declared dep, checks if it already exists at the consuming project's root; if missing, recurses (`incorporate-submodule` on the dep's `ssh_url`).
 - On conflict (same name declared at different ref by two submodules), aborts with a directed error pointing the operator at the conflicting declarations.
 - Emits a final manifest-of-manifests file at `<root>/.helix-manifest.yaml` listing every submodule + its declared deps, for audit + reproducibility.

Anti-bluff guarantee. Every manifest **MUST** be paired with a **verification proof**: a Challenge script (per §11.4.27 + CONST-050(B)) that:

- Bootstraps a throwaway consuming project from scratch in a temp directory.
- Runs `incorporate-submodule` against the manifest under test.
- Verifies the produced submodule layout matches the manifest's declarations (every dep present at its declared layout path; no extras; no missing).
- Runs the submodule's own test suite against the bootstrapped layout; asserts pass.
- Captures wire evidence (per §11.4.2) of every step.

A manifest without this verification proof is a §11.4.31 violation of equal severity to a §11.4 PASS-bluff at the dependency-graph layer.

Cascade requirement. Every owned-by-us submodule **MUST** ship `helix-deps.yaml` at its root, recursively (sub-submodules of submodules ship their own). The constitution submodule itself ships a manifest declaring its own deps (currently empty for the universal Constitution; project consumers may have project-specific extensions). The verifier (`scripts/verify-governance-cascade.sh` or its successor) **MUST** check every owned submodule for manifest presence.

Composition. §11.4.31 directly enables §11.4.28 / CONST-051(C) flat-layout enforcement: nested own-org submodule chains can be mechanically flattened because each submodule declares what it needs, and the incorporator places those deps at the root. The §11.4.28(C) constitution-submodule carve-out does not weaken this: a permitted depth-1 constitution-anchored engine still ships its own `helix-deps.yaml`, so the incorporator reconstructs the graph identically — the manifest is the bridge in both cases. Composes with §1 (manifest schema is itself tested for parse + validation), §3 (submodule-first commit discipline applies to manifest changes too), §11.4.12 (manifest changes regenerate any derived docs/diagrams), §11.4.17 (universal — no project-specific assumptions in the manifest format), §11.4.18 (manifest documented in script-doc + external user guide), §11.4.20 (subagent delegation for cross-submodule manifest authoring), §11.4.25 (manifest presence in coverage ledger), §11.4.26 (constitution-update workflow when extending manifest schema), §11.4.27 (manifest test matrix), §11.4.28 (this rule is its operational complement), §11.4.29 (manifests use snake_case names + canonical paths), §11.4.30 (manifest is tracked source — not a build artefact), CONST-047 (manifests cascade recursively).

Classification: ^{1 1} ⁴ ^{3 2} universal (per §11.4.17). No escape hatch. Severity-equivalent to a §11.4 PASS-bluff at the dependency-graph layer.

2 0 2 6 0 5 1 5

§ . . — Post-Constitution-Pull Validation Mandate (User mandate, — —)

Forensic anchor — verbatim user mandate (2026-05-15):

"Every time we fetch and pull new changes on constitution Submodule we MUST process the whole project and all Submodule (deep recursively) for validation and verification taht every single rule or mandatory constraint is followed and respected! If it is not, IT MUST BE!"

Operative rule. Whenever a consuming project's constitution submodule is fetched + pulled with **any** content change (rule addition, rule revision, gate addition, anchor reference, version bump), the consuming project MUST execute a full-project + recursive-submodule validation sweep BEFORE the new constitution HEAD is treated as canonical for any other work.

Validation sweep contract. The sweep is implemented as `scripts/verify-all-constitution-rules.sh` (canonical name) which:

1. Re-runs the existing governance-cascade verifier (`scripts/ verify-governance-cascade.sh`) covering every §11.9 + CONST-* anchor across every owned submodule (recursive per CONST-047).
2. For each rule whose enforcement gate is implementable programmatically (e.g., CONST-053 `.gitignore` -pattern audit; CONST-051(C) nested-own-org-chain audit; CONST-052 case-conformance audit; CONST-050(A) mock-path-from-production- code audit; CONST-035 anti-bluff smoke-scan), the sweep runs the corresponding gate against the post-pull tree.
3. Any failure produces a directed FAIL entry naming the rule (e.g., `FAIL: CONST-053 – *.log tracked at the-project/foo.log`)
 - the canonical fix.
4. Failures populate the project's Issues tracker per §11.4.15 with Status: `Reopened` (since a previously-passing audit now fails) and Type: `Bug` (since real codebase state violates the constitution).
5. The agent or operator MUST resolve every FAIL before treating the new constitution HEAD as canonical — closure of each reopened item per the existing §11.4 anti-bluff covenant (positive-evidence-only, captured wire evidence).

Pull-time invocation. `git submodule update --remote constitution` MUST trigger the sweep automatically (post-update hook OR the operator's commit wrapper invokes it as part of any commit that advances the constitution submodule pointer). Operator-explicit manual invocation MUST also be available (`./scripts/verify-all-constitution-rules.sh`).

Anti-bluff guarantee. A sweep that exits PASS without actually running every implementable gate is a §11.4.32 violation. The sweep's own meta-test (paired mutation, §1.1) MUST plant a known violation of each enforced gate and assert the sweep reports FAIL for the planted gate.

Composition. §11.4.32 is the **enforcement engine** for every other §11.4.x and CONST-NNN rule. Without it, new rules cascade as anchors but never get enforced in the codebase. Composes with every rule that has a programmatic gate: §1, §1.1, §11.4.10, §11.4.12, §11.4.15, §11.4.16, §11.4.18, §11.4.20, §11.4.22, §11.4.24, §11.4.25, §11.4.26, §11.4.27, §11.4.28, §11.4.29, §11.4.30, §11.4.31, CONST-035, CONST-038, CONST-042, CONST-043, CONST-044, CONST-045, CONST-046, CONST-047, CONST-048, CONST-049, CONST-050, CONST-051, CONST-052, CONST-053, CONST-054.

Classification: universal (per §11.4.17). No escape hatch. Severity-equivalent to a §11.4^{11 4 30} PASS-bluff at the constitutional- enforcement layer — without §11.4.32, every other rule is a decorative anchor rather than an enforced gate.

2026 05 15

§ . . – .gitignore + No-Versioned-Build-Artifacts Mandate (User mandate, – –)

Forensic anchor — verbatim user mandate (2026-05-15):

"every project module, every Submodule, every service and application MUST HAVE proper .gitignore file! We MUST NOT git version build artifacts, cache files, tmp files, main .env file(s) or any files containing sensitive data, API keys or token! Any build derivative which we can recreate by executing proper mechanism for generating MUST NOT be versioned! We MUST pay attention what is going to be committed every time we are preparing to execute commit! If any violation is detected it MUST be fixed before commit is executed!"

Operative rule. Every project module, owned-by-us submodule, service, and application under this Constitution MUST ship a proper `.gitignore` covering its full set of forbidden-from-version- control patterns. The following file/directory classes MUST NEVER appear in version control:

1. **Build artefacts** — anything produced by a build / compile / package step that can be regenerated from sources:

- Binaries (`/bin/` , `/build/` , `/dist/` , `/out/` , `target/` , `*.exe` , `*.dll` , `*.so` , `*.dylib` , `*.a` , `*.o` , `*.class` , `*.pyc`).
- Generated source files where the generator is committed (e.g. protobuf `.pb.go` may be checked in OR ignored — project decides — but never both).
- Bundled assets where the bundler is committed.

2. **Cache files** — anything a tool regenerates on demand:

- `__pycache__/` , `.pytest_cache/` , `.mypy_cache/` , `.ruff_cache/` , `node_modules/` , `.next/` , `.nuxt/` , `.cache/` , `.gradle/` , `.idea/cache/` , `.vscode-test/` , `target/` , `.terraform/` , language-server caches.

3. **Temp files** — `*.tmp` , `*.swp` , `*.swp` , `*~` , `.DS_Store` , `Thumbs.db` , IDE backup files, `*.orig` , `*.rej` .

4. **Sensitive-data files** — anything containing credentials, tokens, keys, or personal data:

- `.env` , `.env.*` (allow `.env.example` / `.env.sample` / `.env.template` with placeholder values only — never real secrets even as examples).
- `*.pem` , `*.key` , `*.crt` , `*.p12` , `*.pfx` , `id_rsa*` , `id_ed25519*` , `*.kdbx` , `.netrc` , `secrets/` directory trees.
- `api_keys.sh` , any file containing `BEGIN PRIVATE KEY` , credential tokens, or session cookies.

5. **Generated reports / logs** — `*.log` , `coverage.out` , `*.coverage` , `htmlcov/` , `*.gcda` , `*.gcno` , screenshots/ recordings that aren't reference assets.

6. **OS / IDE / personal-state files** — `.idea/` , `.vscode/` (except shareable `.vscode/settings.json` etc. if explicitly project-shared), `.history/` , `.svn/` , editor-state files.

7. **Test playback artefacts** (§11.4.14) — every test's runtime capture goes under an ignored evidence directory unless the capture IS the reference asset.

Anti-bluff invariant. Per CONST-035 / §11.4 the presence of a `.gitignore` line alone is not sufficient — the project MUST also verify that no file matching the forbidden patterns is currently tracked. A `.gitignore` that lists `*.log` while `app.log` sits tracked in the repo is a §11.4.30 violation of equal severity to no `.gitignore` at all.

Pre-commit attention. Every commit author (human OR agent) MUST inspect `git diff --staged` and `git status` BEFORE executing the commit. If a staged path matches any forbidden class, the commit MUST be aborted and the issue fixed (un-stage the path, add to `.gitignore` , scrub if already-tracked). Gate `CM-GITIGNORE-PRECOMMIT-AUDIT` : pre-commit hook (or commit wrapper in §2) inspects every staged path against the forbidden-pattern matrix; hits abort the commit with a directed error pointing the operator at the specific offending path and the canonical fix. Paired mutation (§1.1): stage a `*.log` → gate FAILs.

Cascade reach. This rule applies recursively to every owned-by- us submodule per CONST-047 and to every owned-submodule's nested non-owned trees within their working copy. The `.gitignore` files themselves are project-specific in their concrete patterns but universally bound to the rule's normative force.

Secret-leak intersection. §11.4.30 composes tightly with §11.4.10 (Credentials-handling mandate / CONST-042) and §12.1. A `.env` leak is BOTH a §11.4.30 violation (build/sensitive file versioned) AND a §11.4.10 violation (credential leak). The combined severity is **release-blocker requiring rotation + post-mortem** per §11.4.10.

Coverage of "recreatable" content. When in doubt about whether something is a build derivative: if there exists a documented, scripted, or automated mechanism that recreates the file from sources, the file is a build derivative and MUST be ignored. The committed sources MUST include the generator (Makefile target, npm script, codegen invocation), so any downstream consumer regenerates the artefact on demand.

Classification: universal (per §11.4.17). No escape hatch beyond the explicit exceptions enumerated above (e.g., `.env.example` placeholder files). Severity-equivalent to a §11.4 PASS-bluff at the repository-hygiene layer — a tracked build artefact silently drifts vs. its source over time, producing the same class of "works-on-my-machine" surprise that the §11.4 anti-bluff covenant forbids at the test layer.

Composition. §11.4.30 composes with §1 (test coverage — ignored files don't get spuriously included in coverage stats), §2 (commit-wrapper enforces the pre-commit gate), §9.1 (destructive-operation safeguards — never `git clean` an operator's working tree to "fix" a violation), §11.4.10 (credentials-handling), §11.4.12 (auto-generated docs sync — generators committed, outputs ignored only if recreatable on demand from the same generator), §11.4.17 (universal — applies to every consuming project), §11.4.18 (script docs — generator scripts documented), §11.4.20 (subagent delegation for cross-submodule `.gitignore` audit sweeps), §11.4.25 (coverage ledger lists `.gitignore` presence per submodule), §11.4.26 (constitution-update workflow — the constitution submodule's own `.gitignore` complies), §11.4.27 (no-fakes: test fixtures that ARE the reference asset are tracked; runtime captures are ignored), §11.4.28 (submodules-as-equal-codebase: each submodule's `.gitignore` audited on equal basis), §11.4.29 (the renamed paths in the lowercase migration MUST also be properly ignored where applicable), CONST-047 (recursive governance reach^{11 4 29 2026 05 15}).

§ 11.4.29 — Lowercase-Snake_Case-Naming
Mandate (User mandate, 2026-05-15)

Forensic anchor — verbatim user mandate (2026-05-15):

"naming convention for Submodules and directories (applied deep into hierarchy recursively) - all directories and Submodules MUST HAVE lowercase names with space separator between the words of '_' character (snake-case)! All existing Submodules and directories which are not following this rule MUST BE renamed! However, since this will most likely break some of the functionalities renaming we do MUST BE applied to all references to particular Submodule or directory! Everywhere where particular Submodule directory are referenced proper updates MUST BE applied

- all configuration files, documentation and relevant materials, links to Submodules and directories, source code that points to them, etc. There MUST BE reasonable exceptions for this rules - source code for programming languages or Submodules which apply different naming convention - Android, Java, Kotlin and others. Root directory for such applications, services or Submodules can follow OUR convention, but EVERYTHING inside still MUST follow language / technology specific rules! We apply this rules per common sense basis and it MUST NOT be the cause of bigger issues such as technology breaking! Upstreams directory which all of our projects and Submodules have MUST BE renamed to the lowercase letters too, however root project containing the install_upstreams system command (it is exported in our paths in our .bashrc or .zshrc) MUST BE updated to fully work with both Upstreams and upstreams directory. That change if it is not already applied MUST BE done, committed and pushed! ... NOTE: Rules lowercase / snake-case do apply to all project files as well and references to it and from them! Every change done MUST BE covered with all supported test types, full automation tests for validation and verification and fully applied and followed anti-bluff policy and all anti-bluff rules!"

Operative rule. Every directory, submodule, and file under the parent project's working tree MUST use a **lowercase, snake_case** name (ASCII letters / digits / underscores, words separated by _). Existing names that violate the rule (ExampleModule/ , Challenges/ , Containers/ , ExampleAgent/ , HelixQA/ , Security/ , Github-Pages-Website/ , Upstreams/ , Dependencies/ , etc.) MUST be renamed as part of the migration window opened by this clause. Every reference in the codebase MUST be updated atomically with

the rename: configuration files, documentation, user manuals, diagrams, scripts, source-code imports, links, governance files. **Reference drift after a rename is a §11.4.29 violation** of equal severity to the rename itself.

Exceptions (common-sense scope). The rule **MUST NOT** break language-/technology-specific conventions:

- **Programming-language source roots** that mandate a specific case (Java / Kotlin package paths, Android resource folders, Apple framework directories, C# / Swift project layouts) keep their language-mandated names. The submodule's root directory follows our convention; the language-specific subtree inside follows its own.
- **Vendor / upstream submodules** (third-party orgs not in our owned set) keep their upstream-mandated names — we **MUST NOT** rename a third party's repo.
- **Build-tooling artefacts** (`node_modules/` , `__pycache__/` , `.git/` , `target/` , `build/` , `bin/`) keep their tool-mandated names.

When in doubt, the test "does renaming break the technology?" trumps the `snake_case` rule. The §11.4.29 spirit is operator- ergonomics + reference-discoverability, not pedantic uniformity at the cost of technology compatibility.

Upstreams/ → upstreams/ transition. The constitution submodule's installer (`install_upstreams.sh`) — exported on operator paths via `.bashrc` / `.zshrc` — **MUST** support **both** `Upstreams/` and `upstreams/` directory layouts during the migration window, so existing checkouts keep working while consuming projects rename at their own pace. The installer reads whichever directory exists; if both exist the lowercase wins. After every project under this Constitution has migrated, the uppercase fallback **MAY** be retired by a deliberate amendment, but the migration window remains open as long as ANY owned project still ships the uppercase form.

Project-Toolkit Upstreamable submodule synchronisation. The Upstreamable / Project-Toolkit machinery that propagates governance into every consuming project **MUST** be fetched + pulled before any rename batch, and **MUST** itself comply with this rule. Any Upstreamable submodule lacking BOTH-directory support is a release blocker for the rename program.

Test coverage of renames. Every batch of renames MUST ship with: (i) a regression test that verifies every reference to the renamed entity now resolves to the new name (no stale references left); (ii) a full CONST-050(B) test-type matrix run against the post-rename tree; (iii) anti-bluff (CONST-035) wire-evidence captured during the runtime verification. A rename batch without all three is a §11.4.29 violation.

Cascade reach. This rule applies recursively through every owned-by-us submodule layer (per CONST-047) and applies equally to the consuming project and its owned submodules (per CONST-051). Per CONST-051(C) — dependencies at the parent root — the renamed paths MUST be the only canonical location; old-name aliases remain only as transitional symlinks if absolutely required, and those symlinks MUST be removed on next-N-cycle review (project- configurable, recommended ≥ 3 release cycles).

Phased execution. Because the rename touches every reference, the execution MUST be planned as fine-grained phases per the operator's explicit instruction: comprehensive brainstorming, phase-divided plan, fine-grained tasks/subtasks with enormous detail, every change covered by every applicable test type. The phases run in parallel with mainstream work (§11.4.20 subagent delegation is the natural fit for cross-submodule rename sweeps).

Classification: universal (per §11.4.17). No escape hatch beyond the explicit common-sense exceptions enumerated above. Severity-equivalent to a §11.4 PASS-bluff at the reference-integrity layer — a half-completed rename that leaves broken references is worse than no rename at all because it silently breaks consumers.

Composition. §11.4.29 composes with §1 (four-layer floor for every rename batch), §1.1 (paired mutation: rename without updating a reference → gate FAILs), §11.4.12 (regenerate auto- generated docs on rename), §11.4.17 (universal — no project- specific assumptions), §11.4.18 (script-doc-sync after renamed script paths), §11.4.20 (subagent delegation for cross-cutting rename sweeps), §11.4.25 (every renamed path appears in coverage ledger), §11.4.26 (constitution- submodule rename pipeline), §11.4.27 (rename-touch test types apply across the matrix), §11.4.28 (owned submodule renames cascade per CONST-047), CONST-047 (recursive governance reach).

§ . . — Submodules-As-Equal-Codebase + Decoupling + Dependency-Layout Mandate (User mandate, — —)

Forensic anchor — verbatim user mandate (2026-05-15):

"All existing Submodules in the project that we are controlling and belong to some our organizations (vasic-digital, HelixDevelopment, red-elf, ATMOSphere1234321, Bear-Suite, BoatOS123456, Helix-Flow, Helix-Track, Server-Factory — we can ALWAYS check dynamically using GitHub and GitLab CLIs) are equal parts of the project's codebase! We MUST work on that code as much as we do with main project's codebase! All on equal basis! Equally important! We MUST take it into the account, analyze it, extend it, create missing tests, do full testing of it, fill the gaps (if any), fix any issues that we discover or they pop-up, write and extend the documentation, user guides, manulas, diagrams, graphs, SQL definitions, Website(s) and all other relevant materials! We MUST NEVER modify Submodules to bring into them any project specific context since they all MUST BE ALWAYS fully decoupled, project not-aware, fully reusable and modular (by any other project(s)), completely testable! All Submodule dependencies that are used by Submodule MUST BE accessed from the root of the project! We MUST NOT have nested Submodule dependencies but accessing each from proper location from the root of the project — directly from project's root project_name/ submodule_name or some more proper structure project_name/ submodules/submodule_name! This MUST BE heavily enforced and respected so no chaos and mess is created with various dependencies we may have!"

Operative rule. Three cooperating invariants govern every consuming project's relationship with its owned-by-us submodules (those whose upstream `origin` lives under one of the operator- listed orgs — `vasic-digital` , `HelixDevelopment` , `red-elf` , `ATMOSphere1234321` , `Bear-Suite` , `BoatOS123456` , `Helix-Flow` , `Helix-Track` , `Server-Factory` — or any additional org the operator subsequently authorises, the canonical list discoverable at any time via `gh org list` / `glab` / the orgs' public APIs):

(A) Equal-codebase invariant. Every owned-by-us submodule is an **equal part** of the consuming project's codebase. The consuming project's engineering practice — analysis, extension, test creation, gap-filling, bug-fix, documentation (user manual, guides, diagrams, graphs, SQL definitions, website pages, any other authoring surface) — applies to each owned submodule on equal basis. A round of work that improves only the main project while leaving an owned-submodule deficiency unaddressed is a §11.4.28 violation, severity-equivalent to a §11.4 PASS-bluff at the project-scope layer. Coverage ledgers (§11.4.25) **MUST** list every owned submodule as an in-scope target. CONST-047 (recursive cascade) governs the propagation of governance changes; §11.4.28 is the engineering-content counterpart that mandates the same attention to non-governance content.

(B) Decoupling / reusability invariant. Owned submodules **MUST** remain **fully decoupled** from any specific consuming project. No project-specific context, hardcoded paths, hostnames, asset names, naming schemes, or runtime assumptions may be introduced into an owned submodule's source tree. Every owned submodule **MUST** be:

- **Project-not-aware** — its code, tests, and docs make no reference to which parent project consumes it.
- **Fully reusable** — any future Helix-family or third-party project must be able to consume the submodule unmodified.
- **Modular** — its public surface is the only documented integration contract; internal layout may evolve without breaking consumers.
- **Completely testable** — every public surface has standalone tests (per the §11.4.27 100%-test-type matrix) that pass when the submodule is checked out as a standalone repo, without any parent-project rigging.

A commit that adds `project_name/...` strings, hostnames belonging to a specific deployment, or other parent-project context to a submodule's source tree is a §11.4.28 violation. The honest path when a submodule needs information from the parent project is configuration injection (env var, config file, constructor parameter) — never a hardcoded reach into the parent's tree.

(C) Dependency-layout invariant. Every dependency that an owned submodule itself consumes **MUST** be accessible **from the root of the parent project** at one of two canonical paths:


```
<project_root>/<submodule_name>/          # flat layout
<project_root>/submodules/<submodule_name>/  # grouped layout
```

Nested-submodule chains are FORBIDDEN. A submodule MUST NOT have its own `.gitmodules` entries that pull in further owned- by-us repos (transitive own-org submodule recursion). Every dependency required by submodule X MUST be added to the parent project at the canonical path above; X reaches it via documented import / SDK path / runtime resolver — never via its own nested submodule pointer.

EXCEPTION — constitution-submodule-anchored reusable engines (operator Option-2 decision, 2026-07-07; universal per §11.4.17). The single scoped carve-out to the nested-chain prohibition: THE CONSTITUTION SUBMODULE ITSELF (the canonical root per §11.4.35, where reusable engines are anchored) MAY host its own reusable- engine submodules under

`constitution/submodules/<name>/` , PROVIDED each such nested submodule (i) ships a §11.4.31 `helix-deps.yaml` at its root so consumers can reconstruct the dependency graph, AND (ii) declares ZERO further own-org submodules of its own — DEPTH-1 ONLY, never a recursive chain (a nested engine that itself nests an own-org submodule is still FORBIDDEN). This EXCEPTION applies to NO OTHER repository: arbitrary consumer submodules remain forbidden from nesting and MUST use the root-flattened `helix-deps.yaml` bridge above. The rationale for the carve-out is that the constitution submodule is the deliberate home of reusable, project-agnostic engines (§11.4.28(B) / §11.4.74) — anchoring them as depth-1 nested submodules there, each with a manifest, keeps the version-drift risk the base rule guards against bounded (one authoritative home, one manifest, no transitive recursion) while satisfying the operator's intent that the engines be *derived from / anchored by* the constitution.

Rationale: nested own-org submodule chains cause version-drift chaos (two consumers of `LLMsVerifier` end up at different SHAs because each parent picked a different transitive path). The flat / grouped layout makes the consuming project's submodule graph a tree-of-depth-1, which any developer can audit at a glance via `git submodule status` from the project root.

Third-party submodules (not under our orgs) are exempt — they MAY appear at any depth as the upstream's structure dictates. The invariant applies only to our owned set.

Audit + enforcement.

- Gate `CM-OWNED-SUBMODULE-EQUAL-ENGINEERING` (project-side): every release-gate sweep verifies each owned submodule has current test runs, coverage entries (§11.4.25), and documentation freshness on par with the main project. Stale submodules surface as §11.4.28 violations (Status: Operator-blocked or In progress per §11.4.21 — never "ignored").
- Gate `CM-OWNED-SUBMODULE-DECOUPLING` (submodule-side): every owned submodule's pre-commit hook (or equivalent) greps the staged diff for parent-project names / hostnames / asset names. Hits abort the commit until refactored to configuration injection.
- Gate `CM-OWNED-SUBMODULE-LAYOUT` (project-side): the parent project's pre-merge sweep verifies (i) every owned submodule sits at `<root>/<name>/` or `<root>/submodules/<name>/`, (ii) no owned submodule contains a nested `.gitmodules` entry whose upstream is in our org list — EXCEPT the constitution submodule, which MAY carry depth-1 nested own-org engine submodules under `constitution/submodules/<name>/` PROVIDED each ships a `helix-deps.yaml` AND declares zero further own-org submodules (the §11.4.28(C) carve-out; a depth- ≥ 2 chain, or a nested engine missing its manifest, still FAILs), and (iii) every dependency declared by an owned submodule has a corresponding parent-project submodule entry at the canonical path. Paired §1.1 mutation: add a SECOND-level own-org `.gitmodules` entry inside a constitution-anchored engine (a depth-2 chain) → gate FAILs; OR drop the engine's `helix-deps.yaml` → gate FAILs; restore → PASSes.
- Paired mutations (§1.1) for all three gates: plant the forbidden pattern → gate FAILs; restore → gate PASSes.

Workflow integration. Honoring §11.4.28 in practice:

- Engineering rounds plan in submodule-aware tranches: "improve feature X in main; in same round, audit + improve the X- related surface in submodule Y; cascade governance per CONST-047 as usual".

- The §11.4.25 coverage ledger row format is extended with a submodule column so coverage rolls up across the whole owned set.
- Cross-submodule refactors that touch shared types or interfaces ship as a single change-window: parent + every affected owned submodule advanced in lockstep (§3 submodule- first-commit discipline + §2.1 multi-upstream push).

Classification: universal (per §11.4.17). No escape hatch. Composes with: §1 (four-layer test floor reaches submodules too), §1.1 (false-positive immunity), §3 (submodule changes propagate through submodule commits first), §11.4.17 (universal-vs-project classification), §11.4.20 (subagent delegation for the cross-submodule audit sweep), §11.4.25 (full-automation coverage ledger expanded to submodules), §11.4.26 (constitution-update workflow's submodule-pointer-bump step), §11.4.27 (100%-test- type coverage applies to every submodule's standalone surface), CONST-047 (recursive governance cascade). A round of work that violates §11.4.28 is a release blocker for the consuming project, severity-equivalent to a §11.4 PASS-bluff at the codebase-completeness layer.

§ . . — No-Fakes-Beyond-Unit-Tests +
 %-Test-Type-Coverage Mandate (User
 mandate, — —)

Forensic anchor — verbatim user mandate (2026-05-15):

"Mocks, stubs, placeholders, TODOs or FIXMEs are allowed to exist ONLY in Unit tests! All other test types MUST interact with real fully implemented System! No fakes, empty implementations or bluffing is allowed of any kind! All codebase of the project MUST BE 100% covered with every supported test type: unit tests, integration tests, e2e tests, full automation tests, security tests, ddos tests, scaling tests, chaos tests, stress tests, performance tests, benchmarking tests, ui tests, ux tests, Challenges (fully incorporating our Challenges Submodule — <https://github.com/vasic-digital/Challenges>). EVERYTHING MUST BE tested using HelixQA (fully incorporating HelixQA Submodule — <https://github.com/HelixDevelopment/HelixQA>). HelixQA MUST BE used with all possible written tests suites (test banks) for every applications, service, platform, etc and execution of the full HelixQA QA autonomous sessions! All required dependency Submodules MUST BE added into the project as well (fully recursive!!!)."

Operative rule. Two cooperating invariants:

(A) No-fakes-beyond-unit-tests. Mocks, stubs, fakes, placeholders, in-memory facsimile implementations, `TODO`, `FIXME`, "for now", "in production this would", or any empty-implementation pattern are PERMITTED only inside unit-test sources (e.g., `*_test.go` files invoked WITHOUT the integration build tag; `tests/unit/`; equivalent per-language conventions). Every other test type — integration, end-to-end, full automation, security, DDoS, scaling, chaos, stress, performance, benchmarking, UI, UX, Challenges, HelixQA suites — MUST exercise the **real, fully implemented system** against real infrastructure (real databases, real HTTP endpoints, real containers, real downstream services, real captured devices). A non-unit test that imports mocks, in-memory repositories, fabricated provider responses, or placeholder structs is a §11.4.27 violation regardless of how green its summary line looks — severity-equivalent to a §11.4 PASS-bluff. The same prohibition extends to **production code**: no mock import path may be reachable from any code path that runs in production binaries. Pre-build gate `CM-NO-FAKES-BEYOND-UNIT-TESTS` scans the non-unit test trees and production trees for the forbidden patterns; paired mutation plants a fake → gate FAILs.

(B) 100% test-type coverage with every supported type. Every project under this Constitution MUST cover its entire codebase with **every supported test type** the project's domain warrants:

1. **Unit** — fast, isolated, mocks permitted per (A).
2. **Integration** — multi-component, no mocks, real backing services.
3. **End-to-end (E2E)** — full user-flow exercise on target topology.
4. **Full automation** — orchestrated suites exercising every feature × platform combination (§11.4.25 coverage ledger).
5. **Security** — authn/authz boundaries, secret-leak scans (§11.4.10), input-fuzzing, dependency-CVE scanning, threat- model verification.
6. **DDoS** — request-flood resilience at advertised throughput tier, with rate-limit + back-pressure assertions.
7. **Scaling** — horizontal + vertical scale behaviour under linear load growth, including replica add/remove transitions.
8. **Chaos** — controlled failure injection (network partition, process kill, disk full, clock skew) verifying graceful degradation paths.
9. **Stress** — sustained load above advertised tier, asserting bounded resource exhaustion + clean recovery.
10. **Performance** — latency / throughput / tail-latency invariants vs SLO baselines.
11. **Benchmarking** — micro + macro benchmark suites with historical p95-drift detection (§11.4.24 build-resource composition).
12. **UI** — visual-regression + DOM-state + interaction-flow coverage on every target platform's UI surface.
13. **UX** — flow-correctness + accessibility + i18n + visual- cue ordering (§11.4.23 composition).
14. **Challenges** — `vasic-digital/Challenges` submodule fully incorporated; per-feature Challenge scripts covering real user use-cases with captured runtime evidence.
15. **HelixQA** — `HelixDevelopment/HelixQA` submodule fully incorporated; ALL written test banks executed; full autonomous QA sessions run as part of release gates.

Per-type 100% coverage means: for every feature × platform cell in the §11.4.25 coverage ledger, the cell carries a verified PASS evidence pointer for **each supported test type the cell warrants** (a CLI-only feature warrants unit + integration + E2E + full-automation + security + chaos + stress + performance

- Challenges + HelixQA; a UI feature additionally warrants UI + UX; a network service additionally warrants DDoS + scaling + benchmarking). Gaps are tracked per §11.4.15 with explicit UNCONFIRMED: / PENDING_FORENSICS: / OPERATOR-BLOCKED: reasons.

(B.1) Seven-canonical-type enumeration (extension, 2026-08-15). The CANONICAL SEVEN-TYPE ENUMERATION per operator mandate 2026-08-15 (§11.4.224(B.1) extension of the same date) is: (1) UNIT, (2) INTEGRATION, (3) E2E, (4) FULL AUTOMATION, (5) SECURITY, (6) PERFORMANCE AND BENCHMARK, (7) ANTI-BLUFF — the MINIMUM breadth every consuming project satisfies. The broader domain-warranted set enumerated in clause (B) above (DDoS, scaling, chaos, stress, UI, UX, Challenges, HelixQA autonomous QA sessions) is ADDED to the seven-type floor per domain, never SUBTRACTED. Applicability of each type is consumer-owned §11.4.35 DATA — a test type genuinely inapplicable is §11.4.3 SKIP-with-reason at the type-breadth ledger, never silently omitted (§11.4.6 no-guessing at the coverage-claim layer). §11.4.27 sets the type BREADTH; §11.4.224 sets the code-coverage FLOOR across that breadth; both bind together, satisfying one is never evidence for the other (per §11.4.224(F) honest boundary). Propagation: existing CM-COVENANT-114-27-PROPAGATION stays unchanged (literal 11.4.27 unchanged, stays GREEN); CM-NO-FAKES-BEYOND-UNIT-TESTS gate stays unchanged; NEW recommended gate CM-TEST-TYPE-BREADTH-SEVEN-CANONICAL — every consuming project's test suite covers all 7 canonical types where applicable, honest SKIP-with-reason where not; a paired §1.1 mutation that drops one type without a SKIP-with-reason → the gate FAILs.

Required submodule incorporation (recursive). Every project consuming this Constitution MUST add the following as Git submodules (or vendored equivalents authorised by the operator), fully recursively per CONST-047 / §11.4.x cascade:

- Challenges — git@github.com:vasic-digital/Challenges.git
- HelixQA — git@github.com:HelixDevelopment/HelixQA.git

- Any additional functionality submodules under `vasic-digital/*` or `HelixDevelopment/*` orgs that the consuming project depends on (do not duplicate work the organisations already maintain).

Submodule pointers MUST be bumped to upstream HEAD in the SAME commit as any dependent cascade work (§11.4.26 step 7). Pointer drift = §11.4.27 violation.

HelixQA autonomous sessions. "Full HelixQA QA autonomous sessions" means: HelixQA's orchestrator drives the consuming project's running surface end-to-end, executing every test bank the project has registered with it, capturing wire evidence per check (status + body bytes + body-head + duration + screenshots for UI), and producing a session report that is itself committed via the §11.4.22 lightweight doc-sync wrapper. A green autonomous-session report without captured wire evidence is a §11.4 PASS-bluff at the QA-orchestration layer.

Composition. §11.4.27 composes with: §1 (four-layer floor), §7.1 (positive-evidence-only), §11.4.1 (FAIL-bluffs forbidden), §11.4.2 (recorded-evidence), §11.4.3 (per-topology dispatch), §11.4.6 (no-guessing), §11.4.10 (credentials-handling — security test category), §11.4.15 / §11.4.16 (status + type tracking), §11.4.17 (universal-vs-project classification), §11.4.20 (subagent delegation for orchestrating the test-type matrix), §11.4.22 (lightweight doc-sync for coverage ledgers and session reports), §11.4.25 (full-automation-coverage — §11.4.27 is its strict expansion into per-type-of-test territory). It does NOT supersede them. CONST-047 (recursive submodule application) governs the cascade reach.

Classification: universal (per §11.4.17). No escape hatch. A project shipping with mocks reachable from non-unit-test paths OR missing required test-type coverage is **not delivering a stable product** — release blocker for every consuming project, severity-equivalent to a §11.4 PASS-bluff at the release-gate layer.

2026 05 15

§ . . — Type-aware closure-status vocabulary (User mandate, — —)

Forensic anchor — direct user mandate (verbatim, 2026-05-15):

"make sure we use proper wording for the workable item completion status in Issues, Issues_Summary and Fixed docs: - Fixed -> into the proper wording depending on the workable item type (task, feature, and so on) - for example: Fixed, Implemented, Completed. Add this important detail in our root Constitution, CLAUDE.MD and AGENTS.MD so it is ALWAYS respected and followed!"

Classification: §11.4.17-classified **universal** — naming discipline applies to every project that tracks work items by type. Bug closures aren't "implemented", feature deliveries aren't "fixed", and infrastructure refactors aren't either. Type-mismatched closure vocabulary is a quiet but persistent semantic-drift bluff at the documentation layer — it makes greppable releases lie about what shipped.

The mandate. §11.4.15 defined the lifecycle Status closed-set including the closure terminal value `Fixed (→ Fixed.md)`. §11.4.16 defined the Type closed-set `{Bug | Feature | Task}`. §11.4.33 binds the two: the closure terminal value **MUST** agree with the item Type, drawn from this 3-element closed map:

Item	**Type:**	Closure	**Status:**	value
Bug		Fixed	(→ Fixed.md)	
Feature		Implemented	(→ Fixed.md)	
Task		Completed	(→ Fixed.md)	

The `(→ Fixed.md)` suffix is preserved across all three so the existing migration-discipline tooling (`Issues.md → Fixed.md` atomic move per §11.4.19) continues to work without per-Type branching. Generators (`generate_issues_summary.sh` , `generate_fixed_summary.sh` , status-counter helpers, the §11.4.23 colorizer) **MUST** treat the three terminal values as semantically equivalent (all map to "closed, positive evidence captured") while preserving the literal in the emitted document.

No escape hatch. Closing a `Feature` with `Fixed (→ Fixed.md)` or a `Task` with `Implemented (→ Fixed.md)` is a §11.4.33 violation. Pre-build gate (recommended, per consuming project) `CM-CLOSURE- VOCAB-TYPE-AWARE` walks every `Fixed.md` heading + every `Issues.md` heading whose `**Status:**` is one

of the three terminal values, and asserts the Status value matches the item's `**Type:**` per the table. Paired mutation flips a Bug entry's status from `Fixed (→ Fixed.md)` to `Implemented (→ Fixed.md)` → gate FAILs.

Propagation. Composes with §11.4.15 (item-status tracking), §11.4.16 (item-type tracking), §11.4.19 (Fixed-document column alignment), §11.4.23 (colorisation — the closed-state palette applies regardless of which of the three terminal words is used). ³⁴

2026 05 15

§ . . — Reopened-source attribution mandate (User mandate, — —)

Forensic anchor — direct user mandate (verbatim, 2026-05-15):

"when we reopen some workable item (bug, task or feature) we MUST HAVE details on who did reopened this and why (- Reopened status needs: by who, AI or User + details in main Issues doc.). For example, reopened by AI or reopened by real User, why - test failed, manual testing detected problem, etc. Adapt our docs for this - Issues, Issues_SUMmary and Fixed."

Classification: §11.4.17-classified **universal** — every project that uses §11.4.15's `Reopened` lifecycle value benefits from knowing the reopen source + cause. Without this, reopen-thrash patterns (the same item bouncing between Fixed and Reopened across cycles) cannot be diagnosed, and the §11.4.7 demotion-evidence rule loses its forensic counterparty (you cannot evaluate whether a reopen was operator-side observation or agent-side over-restoration without provenance).

The mandate. Every Issues.md (or equivalent project tracker) heading whose `**Status:**` value is `Reopened` MUST carry, within 8 non-blank lines of the heading, a `**Reopened-Details:**` line that captures four sub-facts:

- **By:** `AI` or `User` (the source-of-truth observer who flipped the status). `AI` covers in-loop reopens (test failure, gate regression, captured-evidence retrospect). `User` covers operator- side observations (manual testing, end-user report, design reconsideration).
- **On:** ISO date (`YYYY-MM-DD`).

- **Reason:** one-line cause classification — chosen from a closed vocabulary { test-failed | manual-testing-detected | captured-evidence-contradicts | end-user-report | cycle-re-discovered | design-reconsidered } . Other values are permitted with explicit Reason: <free text> annotation but the closed list MUST be tried first.
- **Evidence:** path to or short description of the captured artefact that justifies the reopen — log file, recording, gate failure ID, operator quote, etc. Reopens without evidence are §11.4.6 / §11.4.7 violations: the reopen IS a demotion-from-Fixed classification change, and demotion requires positive evidence captured under the conditions that re-exposed the defect.

The Issues_Summary.md (or equivalent) Status column MUST distinguish the four Reopened sub-states by source so a sweep query for "reopens by AI in the last 30 days" is mechanically possible. Suggested column rendering: Reopened (AI: test-failed) vs Reopened (User: manual-testing) .

No escape hatch. A Reopened entry without **Reopened-Details:** is a §11.4.34 violation. Pre-build gate (recommended, per consuming project) CM-ITEM-REOPENED-DETAILS mirrors CM-ITEM-OPERATOR-BLOCKED-DETAILS (Phase 39.AT pattern): walks every actionable heading, scans 8 non-blank lines for a **Status:** Reopened line, then continues scanning for **Reopened-Details:** . Missing details line emits WARN initially, hardens to FAIL once backlog is fully populated.

Propagation. Composes with §11.4.6 (no-guessing — the Reason must be drawn from the closed vocabulary or explicitly annotated; no likely/probably/maybe causes), §11.4.7 (demotion-evidence — reopen IS a demotion from Fixed), §11.4.15 (item-status tracking — extends the Reopened value's discipline), §11.4.21 (Operator- blocked discipline — same audit-line pattern with the same gate shape).

2026 05 15

§ . . — Canonical-root inheritance clarity (User mandate, — —)

Forensic anchor — direct user mandate (verbatim, 2026-05-15):

"Parent AGENTS.MD or CLAUDE.MD are located under constitution directory (Submodule) containing these parent (root) which files we are inheriting inside the project! Pay attention to this and make sure you ALWAYS follow this rule! Same applies to the root Constitution file!"

"If needed (and we think it is needed it seems) add into the root Constitution, AGENTS.MD and CLAUDE.MD located in constitution directory (Submodule) which we are inheriting that the root (parent) Constitution, AGENTS.MD and CLAUDE.MD are not the ones in root of the project but inside the constitution directory (Submodule). We are inheriting it and if project specific rules have to be added or project specific constraints which are not universal and reusable, then they go into the Constitution, CLAUDE.MD and AGENTS.MD of the project directory itself!"

Classification: §11.4.17-classified **universal** — every project that consumes this constitution as a Git submodule MUST unambiguously distinguish the **canonical root** (the constitution submodule's three files) from the **consumer extensions** (the project's own three files).

The defect this anchor closes. Loose terminology — "parent CLAUDE.md", "root CLAUDE.md", "main constitution" — used without a file-path anchor causes a quiet but persistent confusion between the inheriting copy and the source-of-truth. AI agents have already been observed editing the consumer-side `CLAUDE.md` when the User's intent was to extend the canonical layer (or vice-versa), causing universal rules to leak into project-specific copies and project-specific rules to be falsely promoted as universal. The defect compounds: once a rule is misfiled, future propagation gates treat the misfile as canonical and silently spread it.

The mandate. The three files in **this constitution submodule** (`constitution/Constitution.md` , `constitution/CLAUDE.md` , `constitution/AGENTS.md`) are the **canonical root** — also called the **parent** files. They contain only universal rules per §11.4.17 — rules that any project consuming this submodule benefits from.

The three files in **the consuming project's repository root** (`<project-root>/CLAUDE.md` , `<project-root>/AGENTS.md` , optionally `<project-root>/Constitution.md` or equivalent) are the **consumer extensions**. They MUST start with the inheritance pointer per the existing `constitution/CLAUDE.md` "How

inheritance works" section. They contain only project-specific rules per §11.4.17 — rules that reference particular hardware, vendor names, regulatory regions, internal asset names, or project-private conventions.

Operative invariants that follow:

1. **When in doubt about which file to edit:** if the rule is reusable across any project, edit the constitution submodule's file. If the rule references project-private specifics (a hardware revision, a vendor SDK version, a regional constraint, a particular service), edit the consumer's file. Default to consumer-side when uncertain (per §11.4.17 — narrower scope is cheap to widen later, the reverse is expensive).
2. **Terminology anchor.** When prose in any file in this constitutional family references "the parent CLAUDE.md" or "the root Constitution," the referent is the constitution-submodule file at `constitution/<filename>` , never the consumer's file. When it references "the project CLAUDE.md" or "this project's AGENTS.md," the referent is the consumer-side file at `<project-root>/<filename>` . AI agents reading this constitution MUST resolve ambiguous pronouns ("the CLAUDE.md", "the constitution") via this rule.
3. **Universal rules added to the consumer side are misfiled.** Per §11.4.17 commit-time check, before adding a `MUST` to the consumer file, the author assesses whether it would benefit other projects that consume this constitution. If yes, lift it to the constitution submodule first.
4. **Project-specific rules added to the constitution submodule are misfiled.** Per §11.4.17, before adding a `MUST` to the constitution submodule, the author assesses whether it references project-private specifics. If yes, demote it to the consumer file (or genericise it to a universal form first).
5. **Propagation gates target both layers but their contents differ.** The existing `CM-COVENANT-114-N-PROPAGATION` gate family verifies that anchor TEXT for §11.4.N exists in EVERY `CLAUDE.md` and `AGENTS.md` across the project. The constitution-submodule files carry the **canonical** text; the consumer files carry **compact propagation pointers** that reference the canonical authority by file path. Both forms count as valid presence; the gate is satisfied when any form is found.

6. **@constitution/CLAUDE.md import (Claude-Code-style) and the pointer-block fallback (Aider/Codex/Gemini-style) are equivalent.** The consumer's CLAUDE.md MUST start with one of:

- The native import: `@constitution/CLAUDE.md`
- The portable pointer-block per the `## INHERITED FROM constitution/CLAUDE.md` heading defined in `constitution/CLAUDE.md` "How inheritance works".

7. **No silent demotion or silent promotion.** Moving a rule between layers MUST be a visible commit — `git mv` of a section if it's a clean clone, or an explicit "Lifted from to constitution per §11.4.35" / "Demoted from constitution to per §11.4.35" line in the commit message. AI agents MUST NOT silently re-author a §11.4.X anchor in the wrong layer and call it propagation.

Pre-build gate (recommended, per consuming project):

- **CM-CANONICAL-ROOT-CLARITY** — verifies (a) consumer's `CLAUDE.md` opens with the inheritance pointer (either `@import` or `## INHERITED FROM constitution/CLAUDE.md` heading), (b) the constitution submodule's three files are present at the expected path, (c) no `## INHERITED FROM` block in the constitution submodule's own files (those ARE the source-of-truth, not consumers).

Propagation. Composes with §11.4.17 (universal-vs-project classification — §11.4.35 defines the file-layer split that §11.4.17 classifies INTO), §11.4.18 (script documentation discipline — same canonical-vs-consumer file-layer rule applies to docs/scripts/). Reading order: this anchor first, then §11.4.17 for the classification rule.

No escape hatch. Misfiling a rule between layers IS a §11.4.35 violation regardless of intent. The fix is `git mv` + commit per invariant 7, not "leave both copies in place to be safe" — duplicates silently diverge over time.

2026 05 15

§ . . — Mandatory `install_upstreams` on
`clone/add` Mandate (User mandate,
— —)

Forensic anchor — verbatim user mandate (2026-05-15):

"Every Submodule or Git repository we add or clone MUST BE upstreams installed using Upstreamable utility which MUST BE available through exported paths of the host system (in .bashrc or .zshrc) using install_upstreams command executed from the root of the cloned (added) repository - only if in it is Upstreams or upstreams directory present with bash script files (recipes) for all repository's upstreams!"

Operative rule. Every time an operator or agent adds or clones a Git repository (`git clone` , `git submodule add` , `incorporate-submodule` per §11.4.31, etc.) under any consuming project of this Constitution, the post-clone procedure MUST be:

1. `cd` into the newly-cloned (or newly-added submodule's) working tree.
2. If the working tree contains an `upstreams/` directory (or legacy `Upstreams/` per the §11.4.29 transition) populated with one or more `*.sh` recipe files declaring upstream Git SSH URLs (the canonical declaration format defined by this constitution submodule's `Upstreams/*.sh` shape), the operator or agent MUST invoke `install_upstreams` from that directory.
3. `install_upstreams` is a host-system utility installed on the operator's `PATH` via `.bashrc` or `.zshrc` export. Implementation lives in the constitution submodule (`install_upstreams.sh`); operators alias / symlink it onto `PATH` once per host setup. The utility reads the recipe files, configures every declared upstream as a named git remote, and fans out `origin` push URLs across all declared upstreams.
4. If no `upstreams/` directory is present, no `install_upstreams` invocation is required.
5. If `upstreams/` is present but `install_upstreams` is not on `PATH` , the operator MUST install it (per the constitution submodule's setup docs) BEFORE making the clone usable.
6. Skipping step 2 when an `upstreams/` directory IS present is a §11.4.36 violation — the next push from that working tree will land on only one upstream, breaking §2.1 (Multi-upstream push is the norm).

Pre-commit attention. Before the first commit lands inside the newly-cloned working tree, the operator MUST verify that `install_upstreams` has been executed (when applicable). The quickest check:

```
git remote -v | grep -c push
```

 reports the expected upstream count. If it reports `1` while `upstreams/*.sh` declares more, abort the commit and run `install_upstreams` first.

Automation. The constitution submodule's tooling (the future `incorporate-submodule` per §11.4.31, the existing `scripts/init-submodules.sh` patterns in consuming projects) MUST auto-invoke `install_upstreams` as part of any clone / add operation when the target tree has an `upstreams/` directory. Operator-explicit manual invocation remains supported.

Gate CM-INSTALL-UPSTREAMS-ON-CLONE . Pre-merge gate inspects every owned-by-us submodule pointer commit and verifies that: (a) if `upstreams/` (or legacy `Upstreams/`) is present in the submodule's tree, (b) the submodule's local checkout has at least as many configured push URLs as the declared recipe count. Paired mutation (§1.1): remove one upstream from local `origin --push` config → gate FAILs.

Composition. §11.4.36 composes with §2 (single-entripoint commit/push wrapper), §2.1 (multi-upstream push is the norm — this rule is its setup-time complement), §3 (submodule changes propagate through submodule commits first — requires multi-upstream parity), §9.2 / CONST-043 (no force-push — the multi-upstream wrap-up makes unforced parity easy), §11.4.17 (universal — every consuming project's clone procedure), §11.4.20 (subagent delegation when batch-cloning), §11.4.28 / CONST-051 (owned-submodule discipline), §11.4.29 / CONST-052 (`Upstreams/` → `upstreams/` transition is exactly this rule's scope), §11.4.30 / CONST-053 (`.gitignore` ignores `upstreams/` only if the project explicitly decides not to track recipes, which is unusual — recipes ARE source of truth and SHOULD be tracked), §11.4.31 / CONST-054 (submodule-dependency-manifest works alongside upstreams recipes: `helix-deps.yaml` lists WHAT to add at parent root, `upstreams/` recipes list WHERE to push the resulting clones).

Classification: universal (per §11.4.17). No escape hatch. Severity-equivalent to §2.1 multi-upstream-push-violation at the clone-time setup layer — without the post-clone install, the working tree silently has a single-upstream blind spot.

§ . . — Fetch-before-edit mandate (User mandate, — —)

Forensic anchor — verbatim user mandate (2026-05-15):

"Make sure that `feedback_fetch_before_edit` memory rule is part of our constitution Submodule - the root Consitution, AGENTS.MD and CLAUDE.MD. Validate and verify that Proejct-Toolkit and all Submodules do inherit all of them! Follow the constitution Submodule documentation for details."

Background. In multi-agent / multi-upstream codebases — where parallel Claude Code, Cursor, Aider, Codex, Gemini CLI, or human operator sessions may operate on overlapping scope — the local working tree's state can lag behind the canonical upstream state by the time any given agent receives a task. Acting on stale local state produces three failure modes documented in the originating session (2026-05-15):

1. **Redundant work** — the agent re-does what a parallel session already finished, wasting tokens and operator review time. (In the originating incident, "Gitee removal" had already been committed upstream by a parallel agent 25 minutes earlier; the second agent's `git remote remove gitee` was a no-op echo of work already done.)
2. **False confidence** — the agent reports completion of work that was already done by someone else, with no mechanical signal that their commit duplicates upstream history.
3. **Divergent history** — if the agent commits a parallel "completion" of already-done work, the result is two siblings of the same change, doubling the conflict surface for the next push attempt to multi-upstream remotes (§2.1).

The mandate. The FIRST git-touching action of any session, on any consuming project that participates in this constitution, MUST be:

```
git fetch --all --prune
git log --online HEAD..@{u}      # parent
git submodule foreach --recursive 'git fetch --all --prune --quiet'
```


If `HEAD..@{u}` is non-empty, the agent MUST integrate (ff-merge, rebase, or — if non-fast-forward — surface to operator per §11.4.4) BEFORE issuing any local edit, scanner run, or test cycle. The fetch step is non-negotiable even when the operator's directive is phrased as "do X immediately" — the 30-second check prevents hour-long conflict reconciliation later.

Scope. Applies to:

1. The consuming project root.
2. Every owned submodule (per §11.4.28) — recursive.
3. The constitution submodule itself (§11.4.26 step 1 makes this explicit for constitution-side edits; §11.4.37 generalises it to ANY edit on the consuming project, not only constitution edits).
4. Any dependency cloned via `incorporate-submodule` (§11.4.31) or `git submodule add` (§11.4.36).

Anti-bluff invariant. The fetch+log check MUST produce captured evidence — the actual `HEAD..@{u}` output, even if empty. Skipping the check on the basis of "I just fetched" or "nothing could have changed in the last N minutes" is a §11.4.6 (no-guessing) violation: the remote state is not knowable without a fetch.

Composition. Composes with §2.1 (multi-upstream push — without the fetch, the agent can't know which upstream has the canonical ref), §11.4.4 (test-interrupt-on-discovery — newly-fetched commits that contradict the planned work are exactly the "freshly discovered defect" that triggers cycle interruption), §11.4.6 (no-guessing — remote state requires fetch, not assumption), §11.4.20 (subagent delegation — parallel subagents MUST coordinate through fetched state, never assumed-local state), §11.4.26 (constitution update workflow — this rule generalises §11.4.26 step 1 to ALL edits, not only constitution-file edits), §11.4.32 (post-constitution-pull validation — fetch-before-edit happens BEFORE the constitution-pull sweep §11.4.32 mandates).

Gates. Pre-build gate `CM-FETCH-BEFORE-EDIT-AUDIT` (when implemented in the consuming project) audits the most-recent commit range against the upstream HEAD at the commit's parent — if the parent ref was not the upstream HEAD at the time the commit was authored, FAIL. Paired mutation (§1.1): synthetic commit whose parent is N commits behind the then-current upstream HEAD — gate must FAIL.

Classification: universal (per §11.4.17). The rule has no project-specific assumptions — it applies to any multi-upstream / multi-agent codebase. No escape hatch. Severity-equivalent to a §11.4 PASS-bluff at the planning layer — acting on stale local state is the operational analog of asserting truth from unverified premises.

§ 11.4.38 – Installable-Asset Evidence Mandate (User mandate, – –)

Forensic anchor — verbatim operator report (2026-05-17):

"app does not have a launcher icon anymore — the latest published app tester build(s). how come app passed anti-bluff checks?"

For any user-distributable build artifact (package, bundle, installer, or container image produced by the build pipeline and distributed to end users), tests and challenges **MUST** open the artifact and verify each user-visible asset is **present** and **non-degenerate**.

"User-visible asset" includes (non-exhaustive):

- Application icons for every density / resolution tier declared in the artifact metadata, including any platform-specific icon format (multi-resolution container, adaptive XML, symbol set, etc.) that the target OS uses when the minimum supported OS version requires it.
- Splash screens declared in the install metadata.
- Application name strings as declared in the install metadata.
- Any other asset whose absence causes the user to be unable to identify, launch, or interact with the installed application from the OS launcher / home screen / app drawer.

A PASS without opening the artifact and verifying the asset chain end-to-end is a §11.4 PASS-bluff, regardless of whether source files exist and tests otherwise pass. The specific failure mode this rule targets is: source file exists → build

pipeline packages it → post-build checks pass at the source layer → artifact ACTUALLY produced with the asset stripped or misconfigured, and no gate ever opens the artifact to verify.

Required evidence: the anti-bluff challenge for a user-distributable artifact MUST produce per-asset PASS/FAIL lines showing: (a) the asset entry exists in the artifact package listing, (b) the asset is non-empty / non-degenerate (size ≥ platform-defined minimum OR format-validated), (c) where the OS's asset-resolution path involves indirection (alias, XML reference chain, density-qualifier override), the full chain is traced to the final rendered resource.

Consuming-project implementation: each consuming project ships one challenge script per artifact type that opens the produced artifact and verifies every declared user-visible asset. The challenge MUST run as part of the project's standard QA gate (equivalent of `make qa-all`).

Root cause of §11.4.38 introduction: a consuming project shipped multiple releases with the application icon absent on the target OS because: (1) the icon XML used a resource type that the OS resolves differently above a specific API level, (2) all pre-ship challenges verified source files only, none opened the packaged artifact. This is a pure §11.4 PASS-bluff — every test and challenge reported green while the end user saw no icon.

Classification: universal (per §11.4.17). No escape hatch. Severity- equivalent to a §11.4 PASS-bluff at the artifact-packaging layer. Composes with §11.4.1–§11.4.5 (evidence requirements), §11.4.25 (full automation coverage), §11.4.27 (no fakes beyond unit tests — source-layer checks of a distributable artifact are the distributable- layer analog of unit-test-only coverage). See Constitution §11.4.38 for the full mandate.

2026 05 17

§ 11.4.38 – Per-Feature On-Device End-User Validation Mandate (iterative, continuous)

Every user-facing feature in a consuming project MUST have at least one HelixQA scenario that exercises the feature from cold-launch with **positive runtime evidence**:

- A screenshot or screen recording captured at a named checkpoint.

- At least one assertion of type `screenshot` , `accessibility_count` , `accessibility_node` , `ocr_text` , `pixel_histogram` , or `editor_state` .
- Evidence files written to a discoverable output directory so every PASS is cross-verifiable after the fact.

Scenarios MUST be RE-EXECUTED on every release candidate to catch cross-iteration regression. A scenario authored but never re-run on subsequent iterations degrades to a metadata-only gate — a §11.4 PASS-bluff at the regression layer.

Authoring rules (consuming project sets project-specific paths):

1. Scenarios are stored in a canonical bank directory designated by the consuming project (e.g. `<banks-root>/feature-coverage/`).
2. Each scenario YAML MUST include: `name` , `version` , `metadata` , `platforms` , `test_cases` , and at least one step with `evidence_required: true` and a specific `evidence_type` .
3. A **coverage matrix** document MUST exist listing feature × iteration × scenario. Matrix rules: (a) new iteration MUST add ≥ 1 scenario per new user-facing feature; (b) prior scenarios MUST NOT be deleted (mark as `status: retired` if feature is removed); (c) all non-retired scenarios MUST PASS for ship-ready.
4. A **static gate** (challenge script) MUST assert the scenario count $\geq N$ (where N is the number of user-facing features registered in the coverage matrix) and that each YAML contains a `evidence_type: assertion`.
5. **Portable host tooling**: evidence-validation scripts MUST use POSIX-portable file-size helpers (e.g. `wc -c < file`) instead of OS-specific commands (e.g. GNU `stat -c%s`) to avoid silent 0-byte reports on BSD/macOS hosts. A regression challenge MUST verify the portable helper on the actual host.

iOS / platform-deferred pattern: scenarios written platform-agnostically with the target platform listed in a comment deferral (e.g.

`# iOS: deferred – tracker #X`) are valid; add the platform to the `platforms:` list when the automation toolchain becomes available — no structural change to the scenario is needed.

Classification: universal (§11.4.17). Composes with §7.1, §11.4 (anti-bluff), §11.4.25 (full automation coverage), §11.4.27 (no fakes), §11.4.38 (artifact evidence). No escape hatch — a feature without an on-device scenario is NOT covered per §11.4.25 invariant 2. See Constitution §11.4.39 for the full mandate.

§ . . — Full-suite retest before
release tag mandate (User mandate,
– –)

Forensic anchor — verbatim user mandate (2026-05-17):

"Please, do complete retests with all existing tests (multi-hours efforts) when all new workable items are done, fixed, polished and verified. Time is essential! We should already have this in our (root) Constitution, CLAUDE.MD and AGENTS.MD."

Operative rule. A release tag (any `vX.Y.Z` / `X.Y.Z-suffix` identifier MUST NOT be created until a **COMPLETE retest with ALL existing tests** has been executed on a clean baseline AFTER every workable item in the batch is done, fixed, polished, and individually verified. A "spot-check" retest that runs only the tests directly touched by the batch is FORBIDDEN — it misses interaction defects between the batch's fixes and previously- stable code paths.

The complete retest comprises:

1. **Pre-build verification full sweep** — every gate in the project's pre-build script runs; 0 FAIL required.
2. **Post-build verification full sweep** — every gate in the project's post-build script runs against the freshly-assembled image; 0 FAIL required.
3. **On-device 4-phase cycle** (e.g. `test_all_fixes.sh` or project equivalent) on **EVERY owned device** (the full topology at that point in time). Each phase (Immediate Fresh Flash / After Reboot / After Factory Reset / After Final Reboot) MUST complete; no phase may be skipped. Phase summaries captured.
4. **Meta-test full mutation sweep** — every paired mutation in `meta_test_false_positive_proof.sh` (or project equivalent) executed; every gate proven non-bluff.
5. **Test bank full sweep** — if a Challenge-driven test bank exists, every Challenge in the bank covered by a full QA session (not a sub-set).

6. **Issues.md / Fixed.md state audit** — no item silently demoted; every Reopened item has §11.4.34 Reopened-Details ; every closure has captured-evidence per §11.4.5; no Status: value outside the §11.4.15 closed-set.
7. **CONTINUATION.md sync check** — per §12.10, document reflects current state at the moment of tagging.

Time is essential. A complete retest is typically a **12–48 hour elapsed effort** depending on parallelism, device count, and meta-test mutation count. This is NOT optional and NOT abbreviated. Operators should plan release cadence with this duration in mind, NOT skip the retest to ship faster. Skipping the retest is the exact "tests passed but feature broken" failure mode §11.4 specifically prohibits.

Composition with §11.4.4. Per-fix retest (`test_all_fixes.sh` run after each individual fix lands) is STILL mandatory per §11.4.4. §11.4.40 is the **additional final integrity check** that runs after the BATCH is complete — catching interaction defects that individual per-fix retests cannot see in isolation.

Composition with §11.4.7. The complete retest is the authoritative captured-evidence baseline for any closure of items present in the batch. If a §11.4.7-promoted item PASSED its per-fix retest but FAILs the full-suite retest, the closure MUST be reverted and the item moved back to `In progress` — captured-evidence-contradicts under same-conditions per §11.4.7.

Composition with §11.4.39. Per-feature on-device end-user validation runs as part of step 3 (on-device cycle) — every feature's wrapper-test fires during the full-suite retest. The two mandates are complementary: §11.4.39 covers individual feature validation breadth; §11.4.40 covers release-time integration depth.

Gate CM-FULL-SUITE-RETEST-MANDATE . Pre-tag gate inspects the release-candidate state and verifies that within the last 72 hours of the candidate commit, evidence exists of (a) full pre-build + post-build run, (b) on-device 4-phase cycle on every owned device, (c) meta-test full sweep, (d) Issues.md/Fixed.md audit pass. Evidence captured in `docs/changelogs/<tag>.md` per §11.4.4(c) requirements. Paired mutation (§1.1): strip `Full-suite retest evidence` block from a changelog → gate FAILs.

Classification: universal (per §11.4.17) — every consuming project's release procedure. No escape hatch — there is no `--skip-full-retest` flag, no `--quick-release` mode. Operators who feel time-pressured to skip the full retest should instead delay the release until the retest completes; shipping unverified code is worse than delayed shipping.

§ 11.4.17 — Pre-Force-Push Merge-First Mandate (User mandate, 2026-05-17)

Forensic anchor — verbatim user mandate (2026-05-17):

"make sure we bring everything from branches to our side before forc push is done! Afer everything is safely and fully merged and all potential conflicts (if any) resolved, then do force push! make sure nothing isnlost, broken or corrupted on bith sides! add these rules in our root Constitution, CLAUDE.MD, AGENTS.MD (constitution Submodule) if itnis not added already! Extremely important rules and mandatory constraints we MUST HAVE and fully respect!"

Operative rule. Any force-push (`git push --force` , `git push --force-with-lease` , `git push +<ref>` , equivalent history-rewriting operation on any remote) authorised under §9.2 / CONST-043 MUST be preceded by a mechanical 4-step merge- first pipeline that brings every remote-side commit into the local tree, resolves every conflict carefully, and verifies nothing is lost or corrupted on EITHER side BEFORE the overwriting push is executed.

The 4-step pipeline (mandatory, in order):

1. Fetch every remote-side reference. Run

`git fetch --all --prune --tags` against every configured remote (origin + every upstream). Capture the output.

2. Integrate every divergent commit locally. Compute `git log --oneline HEAD..<remote>/<branch>` for every remote. For every non-empty range, integrate the remote commits into the local tree via the appropriate strategy:

- `git rebase <remote>/<branch>` when the local divergent work is a strict superset that should land on top, OR

- `git merge <remote>/<branch>` when the local + remote histories are independent additions that both deserve preservation, OR
- **operator-confirmed cherry-pick** when remote contains a subset of commits already present locally in different form.

3. **Audit the integrated tree.** Verify (a) no conflict markers anywhere in the working tree (`grep -rn '^<<<<<< \|^===== $\|^>>>>>> '` . returns empty across all governance + source + test files — third-party docs containing the markers as illustrative content excepted), (b) no file silently dropped (`git diff --stat HEAD@{1} HEAD` shows only expected additions), (c) every previously-passing test still passes on the integrated tree (per §11.4.4 + §11.4.40 baseline), (d) every captured- evidence artifact from the pre-merge state still validates.
4. **Execute force-push.** Only after steps 1-3 produce captured evidence of clean integration: `git push --force-with-lease <remote> <ref>` (NEVER `--force` without `--with-lease` unless explicitly authorised per §9.2 sub-clause 6 for a specific remote where lease semantics are unavailable). One force-push event per CONST-043 authorisation — no batch authorisation across multiple force-pushes.

Three failure modes prevented:

- (a) **Remote-side content loss.** Force-push without prior fetch silently overwrites commits a parallel session, sibling agent, or co-developer landed on the remote. The lost work is recoverable from the remote's reflog only within its TTL window — typically 30-90 days — but the audit trail (PR comments, CI evidence, governance signatures) is lost permanently.
- (b) **Stale-state act on remote.** Force-push from a divergent local state that was branched from a remote tip N hours ago treats the remote-side N hours of work as nonexistent. The operator's `--force-with-lease` safety check catches the literal SHA mismatch but ONLY when invoked AFTER a fetch that updated the lease reference; without step 1, the lease check is reading stale local refs.
- (c) **Conflict-driven corruption.** A merge that completes with unresolved markers in the tree (one half kept, other dropped silently because the marker was inside a comment block, OR markers committed verbatim as in CONST-049-prerequisite

incidents observed 2026-05-17 on helix_qa + containers) lands broken governance / source on the force-push target. Step 3's explicit conflict-marker grep is the mechanical guard.

Two-gate composition with CONST-043. §11.4.41 does NOT relax CONST-043's operator-approval requirement — it adds a SECOND gate on top:

- **Gate A (CONST-043).** Operator types explicit per-operation force-push authorisation in the conversation.
- **Gate B (§11.4.41).** Agent executes the 4-step merge-first pipeline, captures evidence of clean integration, presents evidence to operator BEFORE the force-push.

Both gates required. CONST-043 alone authorises a force-push that loses remote work; §11.4.41 alone risks force-pushing without operator awareness. The two together produce a force-push that is BOTH operator-authorised AND remote-safe.

Verification artefact. Every §11.4.41-governed force-push emits a `docs/ changelogs/<tag>.md` "Force-push merge-first audit" section containing: (i) `git fetch` output, (ii) `per- remote HEAD..<remote>/<branch>` log before integration, (iii) integration strategy chosen per remote with rationale, (iv) post-integration conflict-marker scan output (must be empty), (v) post-integration test suite delta (must show only expected changes), (vi) the `--force-with-lease` push output with lease SHA evidence, (vii) CONST-043 authorisation quote from the conversation.

Cascade requirement (per CONST-047). This anchor (verbatim or by `§11.4.41 / CONST-061` ID reference) MUST appear in every owned submodule's `CONSTITUTION.md`, `CLAUDE.md`, and `AGENTS.md`. Severity-equivalent to a §11.4 PASS-bluff at the remote-data-integrity layer. Gate `CM-FORCE-PUSH-MERGE-FIRST` walks `docs/changelogs/<tag>.md` "Force-push" entries for the 7 audit elements; paired mutation strips any element and asserts gate FAILs.

Classification: universal (per §11.4.17) — every owned project executes force-pushes against its remotes; the merge- first discipline is reusable verbatim. No escape hatch — the operator-pressure escape ("just force-push, we'll fix it later") is the exact failure mode this anchor closes.

Composes with: §9.2 (data-safety hardlinked backup), §11.4.4 (test-interrupt-on-discovery — broken integration triggers rollback), §11.4.6 (no-guessing — every step's outcome captured, not assumed), §11.4.26 (constitution-submodule update pipeline — that mandate is the per-submodule specialisation of §11.4.41 for governance-files-only commits), §11.4.32 (post-pull validation — the audit step's mechanical companion), §11.4.37 (fetch-before- edit — step 1 enforces it for the force-push specifically), §11.4.40 (full-suite retest — step 3's test-evidence prerequisite), CONST-043 (per-operation operator approval), CONST-047 (recursive cascade).

§ . . – Iteration-discipline mandate (User mandate, – –)

Forensic anchor — verbatim user mandate (2026-05-18):

"Are we clear about working iterations and sorting priorities for testing and fixing? We first fix top and middle critical priorities → We flash and test all of these → If nothing is broken and no new issues are reported by users or found we run full system testing (many hours for everything) → if anything is still broken or new issues are reported we get back to fixing → Cycle repeats. Make sure this is absolutely clear and mandatory foundation of our root (constitution Submodule) Constitution, AGENTS.MD and CLAUDE.MD! We MUST WORK in this manner until otherwise has been told."

Operative rule. Project work proceeds in priority-ordered iteration cycles. Each cycle is a closed loop of five steps that MUST run in order; advancing past any step before its acceptance condition is met is a §11.4 bluff equivalent to skipping the step entirely.

The five-step cycle:

1. **Priority-ordered fix selection.** Open the project's issue tracker, filter to items whose ****Status:**** is one of **Queued | In progress | Reopened | Operator-blocked** (per §11.4.15 closed-set), sort by severity DESC then intra-group criticality DESC (per §11.4.12 sort order), select ONLY items classified **Top critical** or **Middle critical** for the current batch. **Low** items are explicitly DEFERRED until the critical batch is closed.

2. **Batch implementation.** Land source-side fixes for the selected batch per §11.4.9 (batch-source-fixes-before-rebuild) — every non-runtime-dependent fix lands BEFORE the next rebuild, every fix gets the §11.4.4 four-layer coverage (pre-build gate + post-build gate + on-device test + paired mutation).
3. **Smoke-test gate.** After flash, run the SMOKE bundle: (a) anti-bluff baseline (`meta_test_false_positive_proof.sh` for gates touched by this batch), (b) every `test_*.sh` whose name matches a fix in this batch, (c) the critical-path regression probe (boot + launcher + WiFi + audio + secondary-display routing core), (d) the §11.4.39 per-feature on-device validation for any feature touched. Estimated runtime: <30 minutes on the parallel device fleet. Smoke MUST emit zero FAIL and zero unclassified WARN.
4. **Full-system-test gate.** ONLY if (a) smoke is GREEN AND (b) no new operator/user report has surfaced during the batch's validation window, run the §11.4.40 complete retest (all 7 steps including 12–48 h on-device 4-phase cycle on every owned device). If smoke FAILED OR a new report arrived, the full-system test is FORBIDDEN; loop back to step 1 with the new evidence added to the priority queue.
5. **Release-readiness or loop.** If full-system test is GREEN AND no new report arrived during it, the batch is release-ready (the §11.4.40 tagging procedure may proceed pending operator authorization). Otherwise loop back to step 1.

Priority taxonomy (binds existing severity convention). `Top critical` = severity `C` (Critical per Issues.md convention) AND intra-group criticality `5`. `Middle critical` = severity `C` with intra-group `1 – 4` OR severity `M` (Major) with any intra-group score. `Low` = severity `L` per the existing `[C/M/L]` taxonomy documented in §11.4.12. Severity `WARN` items are treated as `Low` for iteration-discipline purposes.

Cycle repeats. The five-step loop continues, fed by the priority queue, until the operator explicitly authorizes a release. There is no upper bound on iteration count; "we've cycled enough" is not a Constitution-recognised stopping condition. Time is essential per §11.4.40 BUT the cycle MUST NOT be truncated to ship faster — that is exactly the bluff §11.4 forbids.

Composition. §11.4.4 (per-fix retest inside step 2) + §11.4.7 (closures require same-conditions positive evidence; step 4 is authoritative baseline) + §11.4.9 (source-side batching inside step 2) + §11.4.34 (Reopened items attribute By: AI / User) + §11.4.40 (the multi-hour full-suite retest IS step 4 — §11.4.42 is the meta-loop conductor binding the instruments).

Anti-bluff coupling. Per §11.4.2 + §11.4.5, every smoke-test and full-system-test PASS MUST carry positive captured evidence of user-visible behaviour. Tests AND HelixQA Challenges bound equally — a Challenge that scores PASS without applicable analysis is a §11.4 PASS-bluff.

No escape hatch. No `--skip-priority-batch` , `--skip-smoke` , `--full-suite-only` , or `--release-without-loop` flag exists. Subagents performing autonomous work default to the §11.4.42 path. Operators who feel time-pressured to short-circuit the loop should add capacity (more devices in parallel, more agent slots) rather than skip steps. Shipping under-validated code is worse than delayed shipping.

Gate `CM-COVENANT-114-42-PROPAGATION` mirrors `CM-COVENANT-114-40-PROPAGATION` 1:1 — every `CLAUDE.md` / `AGENTS.md` in the covenant file set carries the §11.4.42 anchor. Paired mutation strips the anchor literal from one consumer file → gate FAILs.

Classification: ¹¹ ⁴ ⁴³ universal (per §11.4.17) — every consuming project's iteration workflow.
_{2026 05 18}

§ . . — TDD-Fix-Discipline Mandate
(User mandate, - -)

Forensic anchor — verbatim user mandate (2026-05-18):

"Make sure you do validate and verify every fix! Make sure you first do the fix on live device using ADB if that is possible, once fix is confirmed, make sure that it is fully applied into our codebase ... It is important to note that we shall start by creating test which confirms the issue (fails as expected), and ones the fix is produced it becomes positive - it does pass with success! No bluff policy is mandatory and nothing can be bluffed! Every step of this process MUST BE 100% bluff-free and bluff-proofed!"

Operative rule. Every fix MUST follow the 5-step TDD-fix workflow:

1. **RED — Failing test FIRST.** Before any code change, author a test (or extend an existing one) that exercises the user-visible path of the reported defect and FAILs for that defect specifically. Per §11.4.1 the FAIL MUST be a real product defect (not a script-bug induced FAIL). Per §11.4.2 the FAIL run MUST produce captured evidence (logcat / dumsys / screenrecord / dmesg / `/sys` snapshot / sink-side probe). The test identifier MUST cite the §-letter or Fix-# under repair. A RED step that "happens to pass because the underlying probe lifecycle isn't wired" is itself a §11.4 PASS-bluff.

2. **LIVE-ADB-PROBE — try the fix on the running device first (when feasible).**

For mutable surfaces — `setprop persist.*` , `settings put` , `pm clear` , `am force-stop` , `am start` , `cmd wifi` , `cmd media_session` , `ad-hoc / sys` writes, boot- script edits pushed to `/vendor/bin/` , test-fixture replacement — the operator MUST attempt the fix on the running device via `adb shell` first. The fast adb-loop (~5 min per attempt) replaces the build-flash-test loop (~6 h per attempt) for confirmation of the hypothesis. The live-probe IS NOT the fix; it is the EXPERIMENT that proves the hypothesis is correct before the source-side investment.

Live-probe is INFEASIBLE (exception, must rebuild) for: kernel changes, AOSP framework (frameworks/base/), hardware HAL changes, vendor/ partition contents, init.rc service definitions, sepolicy, ro.* build properties (immutable post-boot), AOSP-build-time XML / resource overlays, Android.bp/Android.mk changes, signed-system-app contents (`LOCAL_CERTIFICATE := platform`). Each exception MUST be cited in the commit message as

`LIVE_PROBE_INFEASIBLE: <reason>` .

3. **GREEN — apply the fix to source code.** The source patch MUST achieve the same effect on the device as the live probe. Per §11.4.9 the source change is batched with other source-side fixes where appropriate. Build via the project's containerized build (§12.9). Flash via the project's flash script. The bytes on the assembled image MUST land where the post-build gate expects (per §11.4.4(b) four-layer requirement).
4. **VERIFY — re-run the RED test, must now PASS.** Per §11.4.7 demotion-evidence: the PASS MUST come from the SAME conditions that initially FAILED (same device, same firmware NOW carrying the fix, same load profile, same cycle position). Per §11.4.5 the PASS MUST carry positive captured evidence (presence + correctness — RMS amplitude, fprobe channel count, frame count, OCR text, sink codec state, etc.). Per §11.4.2 video tests MUST cross-check via the recording-analyzer. Per §11.4.42 a reliability check at the iteration count (typically 10 iterations) MUST PASS all iterations — a single intermittent FAIL keeps the item in `In progress` per §11.4.7.
5. **DOCUMENT — update every relevant doc in the SAME commit.** Per §11.4.4(c): `docs/Issues.md` → `docs/Fixed.md` migration with type-aware closure vocabulary per §11.4.33 (Bug → `Fixed` , Feature → `Implemented` , Task → `Completed`); the project's CLAUDE.md Applied Fixes Reference row; per-version `docs/changelogs/<tag>.md` entry; affected user-facing guides under `docs/guides/` ; HelixQA bank entry per §11.4.4(b); `docs/CONTINUATION.md` per §12.10; project memory file if non-obvious.

Composition. §11.4.43 is the workflow that ENACTS the existing covenant.

Composes with: §11.4.1 (the RED FAIL is a real defect), §11.4.2 (captured evidence at both RED and VERIFY), §11.4.3 (topology dispatch in the RED test), §11.4.4 (four-layer coverage at GREEN), §11.4.5 (quality analysis at VERIFY), §11.4.6 (no guessing during the LIVE-PROBE hypothesis statement — `UNCONFIRMED`: until probe proves it), §11.4.7 (demotion-evidence at VERIFY — same conditions), §11.4.8 (research citation precedes step 1), §11.4.9 (GREEN is batched), §11.4.40 (full-suite retest is the release-tag-time additional check), §11.4.42 (10-iteration reliability loop IS step 4's inner cycle).

Gate `CM-COVENANT-114-43-PROPAGATION` . Pre-build gate verifies the §11.4.43 anchor is present in every CLAUDE.md / AGENTS.md across the covenant file set. Paired mutation strips the anchor literal from one consumer file → gate FAILs.

Classification: universal (per §11.4.17) — every consuming project's fix workflow.
No escape hatch. No `--skip-red-test` , `--no-live-probe` , `--skip-verify` ,
`--skip-document` flag exists. "I'll add the test after the fix" is the exact PASS-bluff
pattern §11.4 forbids because the test authored after the fix demonstrates only that
the test agrees with the fix, not that the test catches the bug.

§ . . — Document Revision Header Mandate (User mandate, — —)

Forensic anchor — verbatim user mandate (2026-05-18):

"Add this change / improvement as background work in parallel with our mainstream work. All documents (Issues, Issues_Summary, Fixed, Continuation doc and all others MUST HAVE revision number! Revision number with information such as: last date and time when document was modified MUST GO at the top of the document below the main H1 title. Add all relevant information about the document as well. However, last date and time document has been modified and revision number are MANDATORY! Revision number shall be the whole number 1, 2, 3, etc. Document all this in our root (constitution Submodule) Constitution, AGENTS.MD and CLAUDE.MD as mandatory rule(s) and constraints we MUST HAVE! No bluff is allowed of any kind!"

The mandate. Every IN-scope tracked Markdown document MUST carry a header block directly below the H1 title containing two MANDATORY fields:

- `**Revision:**` N where N is a monotonic positive integer (1, 2, 3, ...). Never decremented. Never reset on rewrite. A freshly- created IN-scope document starts at Revision 1.
- `**Last modified:**` YYYY-MM-DDTHH:MM:SSZ (ISO 8601 UTC).

Optional encouraged fields: `**Description:**` , `**Authority:**` ,
`**Maintainer:**` , `**Scope:**` , `**Auto-generated-from:**` (auto-generated docs only), `**Status:**` (plan docs only).

Header format (canonical):


```
# Document Title
```

```
**Revision:** 47
```

```
**Last modified:** 2026-05-18T13:42:00Z
```

```
**Description:** Open / in-flight bug + work tracker
```

```
**Authority:** Constitution §11.4 covenant
```

```
**Maintainer:** Operator + AI loop per §11.4.42
```

```
<content...>
```

YAML front-matter and HTML-comment headers are FORBIDDEN — they are invisible to humans scrolling the raw Markdown and break the "all relevant information about the document" visibility clause.

IN scope (revision header MANDATORY): `docs/Issues.md` , `docs/Issues_Summary.md` , `docs/Fixed.md` , `docs/Fixed_Summary.md` , `docs/CONTINUATION.md` , `docs/guides/**/*.*md` , `docs/research/**/*.*md` , `docs/scripts/**/*.*md` , `docs/changelogs/**/*.*md` , `docs/superpowers/plans/**/*.*md` , `docs/hardware/**/*.*md` , and every other tracked Markdown under `docs/` .

OUT scope (revision header NOT required): `CLAUDE.md` / `AGENTS.md` (every layer — already version-tracked via `VERSION` file + inheritance pointer per §11.4.35; adding a per-file revision would duplicate authority and obscure the canonical-root contract); `README.md` / `CONTRIBUTING.md` / `LICENSE` / `NOTICE` / `VERSION` / `OWNERS` (standard metadata); rendered HTML/PDF artifacts (revision inherited from source Markdown); auto-generated docs that already carry an embedded generation-time stamp.

Auto-bump tooling:

- `scripts/doc_revision_bump.sh <file>` — manual entry point, idempotent (no-op if Last-modified within 5 s of `date -u`).
- `.git/hooks/pre-commit` — automatic entry point; walks staged IN-scope docs, calls the bump for each, re-stages.
- `scripts/testing/sync_issues_docs.sh` — composes the bump for `Issues_Summary` + `Fixed_Summary` after regeneration.
- `scripts/commit_docs.sh` — calls the bump before staging so operators who skipped the hook installation are still covered.

Composition with §12.10. CONTINUATION.md already carries `Last updated:` per §12.10. §11.4.44 adds `**Revision:**` N directly under H1 + reuses the existing `Last updated:` line as the §11.4.44 `Last modified:` line. One source of truth, two pointers — no duplication.

Composition with §11.4.12. Issues_Summary.md / Fixed_Summary.md inherit the revision number of their source-of-truth Markdown at generation time; they are NOT incremented independently.

Composition with §11.4.23. HTML colorizer preserves the revision header in the colored HTML — pandoc renders the bold lines as styled `<p><strong>...</strong></p>` blocks which the colorizer leaves alone (only Issues-table rows are mutated).

Gates: `CM-DOC-REVISION-HEADER-PRESENT` walks IN-scope docs and asserts both mandatory fields appear within 15 lines below the H1 title; `CM-COVENANT-114-44-PROPAGATION` asserts the anchor literal is present in every covenant file. Paired mutations strip the `**Revision:**` line OR the anchor literal → gates FAIL.

Classification: universal (per §11.4.17). **No escape hatch** — no `--skip-revision-bump`, `--no-header`, `--allow-missing-revision` flag exists. Operators who feel pressured to skip the bump should land smaller commits, not bypass the discipline.

2026 05 18

§ . . – Integration-Status-Doc
Maintenance Mandate (User mandate,
– –)

Forensic anchor — verbatim user mandate (2026-05-18):

"Make sure this Markdown document is regularly updated to reflect current up to date state of full integration ... Whole integrations MUST be regularly updated and retested! Keep the Markdown document ALWAYS up to date and ALWAYS exported into PDF and HTML! Make sure document is ALWAYS in sync! Add all this into our Constitution, AGENTS.MD and CLAUDE.MD."

Generalisation note. §11.4.45 is the generic form of §12.10 (CONTINUATION.md) applied to every domain integration. §12.10 binds the canonical handoff document specifically; §11.4.45 generalises the sync-and-evidence pattern to every non-trivial integration (Dolby/Arvus, NanoKVM, Firebase, Widevine L1/L3, Play Protect, Netflix, Cuttlefish, Chromecast, WiFi chip, ES8388 codec, Sonos, AC-3, Google Stack, HiFi audio, TV apps, etc.). Without this generalisation each new integration re-invents the same sync wrapper, revision-header discipline, captured-evidence requirement, and operator-blocked surface — duplicating effort and risking drift.

Operative rule (10 sub-points — every Status.md MUST hold ALL).

Every integration-status document at `docs/<domain>/<integration>/Status.md` :

1. MUST exist when a domain integration is non-trivial (more than one fix / test / gate landing for that integration).
2. MUST carry the §11.4.44 revision header directly below the H1 title (Revision + Last modified + Description + Authority + Maintainer + Scope).
3. MUST be auto-synced (HTML + PDF) on every related-test-cycle completion AND every fix touching the integration. "Auto-synced" means the wrapper from sub-point 5 is invoked; manual regeneration is forbidden as the sole path because operators will forget.
4. MUST be auto-colored per §11.4.23 — Status column cells colored by lifecycle state, Type cells colored by Type, operator-blocked rows visually distinct.
5. MUST have a sync wrapper invocable as `bash scripts/testing/sync_integration_status.sh <Status.md>` OR a per-integration thin shell that delegates to the generic wrapper. The thin shell exists so operators can type a short command per integration without remembering paths.
6. MUST live under `docs/<domain>/<integration>/` directory structure — consistent layout, discovery by glob remains O(1).
7. MUST include a captured-evidence-driven status table per §11.4.5 — every status claim cites the test log path, recording file path, or sink-probe report that backs it. Status claims without evidence are §11.4 PASS-bluffs.
8. MUST distinguish status values per §11.4.6 / §11.4.7: PASS / FAIL / SKIP / PENDING_FORENSICS / OPERATOR-BLOCKED — closed vocabulary, mechanically distinguishable, no implicit values.

9. MUST list operator-blocked items at the top of the document (just below the revision header) so an operator scanning the Status.md finds action items in O(1) — not buried 200 lines deep in a chronological evidence log.
10. MUST be referenced from docs/CONTINUATION.md §3 (Active work) when any item in the Status.md is non-terminal (Queued / In progress / Reopened / Operator-blocked) — composes with §12.10 so the canonical handoff document points at the integration-status doc rather than re-stating its contents.

Gates:

- CM-COVENANT-114-45-PROPAGATION — anchor text propagated to every CLAUDE.md / AGENTS.md across the covenant file set.
- CM-AF-INTEGRATION-STATUS-DOCS — discovers all docs/**/Status.md via glob, verifies each carries the §11.4.44 header, has a sync wrapper, HTML+PDF exports are mtime-current vs the source, colorization applied (cell-status class present in the colorized HTML).

Paired meta-test mutations (§1.1 compliance): propagation strip → propagation gate FAILs; delete `**Revision:**` from a Status.md → status-docs gate FAILs; touch Status.md without re-syncing → mtime mismatch FAILs; inject `likely` outside `UNCONFIRMED:` block → existing `CM-NO-GUESSING-MANDATE` catches it (composition catch).

Composition. §11.4.5 (captured evidence in every row), §11.4.6 / §11.4.7 (status vocabulary closed-set), §11.4.12 (sync wrapper pattern reused for HTML+PDF re-export), §11.4.13 (sink-side captured evidence is a specific instance of the generic rule), §11.4.15 (status values), §11.4.22 (commit_docs.sh wrapper), §11.4.23 (colorizer extends to Status.md HTML), §11.4.44 (every Status.md carries the revision header), §12.10 (CONTINUATION.md references Status.md paths in §3).

No escape hatch — no `--skip-status-sync`, `--no-revision-bump-on-status`, `--allow-stale-html` flag. Operators who feel pressured to skip the sync should land smaller, more frequent integration commits rather than batching status updates into one giant push.

Classification: universal (per §11.4.17).

§ . . – Validate-recent-work-before-post-flash-tests mandate (User mandate, – –)

Forensic anchor — verbatim user mandate (2026-05-18):

"We see that on newly flashed devices we execute all post flash tests. Make sure we do execute them all after we validate and verify all recent work from the Issues and Continuation docs first! Once all new / latest work is fully validated and verified, no new issues found, nothing is discovered broken or faulty, then we run all post-flash tests! If validation and verification of recent work fails for any reason, and we must go back to fix anything that popped up we do not have to run post-flash tests! We MUST save time! All post-flash tests can be executed only after recent work on features and issues is fully validated and confirmed with complete anti-bluff policy enforced! Add this so it is absolutely clear into our root (constitution Submodule) Constitution, AGENTS.MD and CLAUDE.MD! Start following all these rules IMMEDIATELY! Iff required abort all post-flash testing if it is in progress!"

Operative rule. After every device flash, the orchestrator MUST first run a recent-work validation pass before any full post-flash suite (`test_all_fixes.sh` or project equivalent). The validation pass enumerates recent work from three docs (Issues.md / Fixed.md / CONTINUATION.md §3), maps each recent item to its on-device test binding, runs ONLY those tests, classifies each result per §11.4.6, and demands 100% green before the full suite is permitted.

Definition of "recent work" (closed-set):

- `docs/Issues.md` actionable headings with `**Status:**` ∈ {In progress, Ready for testing, Reopened} (per §11.4.15 closed-set).
- `docs/Fixed.md` headings closed within `--max-age-days N` (default 7).
- `docs/CONTINUATION.md` §3 "Active work" sub-headings listed as IN PROGRESS or BLOCKED.

Recent-work test binding. Each recent item's on-device test is discovered via §-letter or Fix # match against the project's test_*.sh filenames OR an explicit Test: line in the item heading. Items without an on-device test are §11.4.4 four-layer-coverage violations (a fix without an on-device test is forbidden) — flagged as VALIDATION-PASS-BLOCKED until the test lands.

Authorization sequence (machine-enforced):

1. scripts/testing/recent_work_validate.sh --device <serial> runs; writes /data/local/tmp/.recent_work_validated on the device IFF all targeted tests return PASS with captured evidence.
2. The full suite (test_all_fixes.sh) checks for the marker file at entry; absent or stale (older than the last flash boot epoch) → refuses to proceed with exit code 11.
3. Marker is invalidated automatically on reboot — the orchestrator stores the device boot epoch in the marker and re-validates.

Anti-bluff during validation. Each recent-work item that landed a fix in this batch MUST have a paired §11.4.43 RED test (captured before the fix) AND the same test now GREEN (captured after the fix). A GREEN with no prior RED is itself the bluff §11.4 forbids — the gate FAILs.

Honest classification of validation FAILs. Per §11.4.6: PRODUCT defect (real regression — block full-suite, surface to operator), TEST-INFRA defect (test-script bug — fix per §11.4.1 at source layer, re-run), BLUFFY-THRESHOLD (e.g. SLO floor wrong — re-baseline), or UNCONFIRMED (not reproducible in isolation — tag per §11.4.7 PENDING_CYCLE_RETEST, NEVER silently demote).

Composition. §11.4.4 (STOP-on-discovery) + §11.4.6 (no- guessing) + §11.4.7 (demotion-evidence) + §11.4.40 (full-suite gate) + §11.4.42 (iteration discipline — smoke-test ≡ recent-work- validation) + §11.4.43 (RED test for each recent-work item IS the validation test) + §11.4.44 (revision header determines recency) + §12.10 (CONTINUATION.md source-of-truth for §3 Active work).

Gates:

- CM-COVENANT-114-46-PROPAGATION — anchor literal present in every CLAUDE.md / AGENTS.md across the covenant file set.

- `CM-AF-RECENT-WORK-VALIDATION-GATE` — asserts `scripts/testing/recent_work_validate.sh` exists + executable + sources the anti-bluff library + the full-suite orchestrator contains the entry-gate check.
- `CM-AF-VALIDATION-ARTIFACT-FILE` — asserts the marker file path `/data/local/tmp/.recent_work_validated` is literal-identical across helper + orchestrator + propagation-block + this Constitution section.

Paired mutations strip the anchor / remove the entry-gate / flip the marker path → respective gate FAILs.

No escape hatch. No `--skip-validation`, `--full-suite-always`, `--ignore-recent-work` flag is permitted. The full-suite-without-validation pattern is the precise "tests passed but feature broken" failure mode §11.4 specifically prohibits. The 60–90 min validation pass is ALWAYS faster than 4–6 h of full-suite chasing a defect.

Classification: universal^{2026 05 18} (per §11.4.17).

§ . . — Firebase Data Review Mandate (User mandate, — —)

Forensic anchor — verbatim user mandate (2026-05-18):

"We MUST regularly before every bigger working round check Firebase information gathered so far: Crashlytics (all fatals, non-fatals and ANRs), Analytics data and Performance data. For anything problematic depending on severity proper workable items (issues) MUST be created (Issues, Issues_Summary docs) with full references (links) to original data on Firebase. We MUST make sure we do not create for same issues multiple entries, so we MUST distinguish between same issue manifested in different forms (stacktraces)! Everything MUST BE comprehensive and in-depth level of details so any changes we do (or fixes) do solve the original root problems!"

The mandate. Before every "bigger working round" (pre-build, pre-flash, pre-tag, daily, post-deployment burn-in) the operator/ loop MUST execute the Firebase review pass via `scripts/firebase/review_round.sh`. The pass queries

Crashlytics (fatals + non-fatals + ANRs) + Analytics + Performance, classifies each finding by the §11.4.47 severity table, dedup-maps to existing Issues.md entries via the three-tier algorithm, and drafts new Issue entries for unrecognised findings. Skipping the pass is a §11.4 PASS-bluff: Firebase IS the captured evidence from real end-user devices; ignoring it is the precise "tests pass, feature broken in the wild" failure mode §11.4 prohibits.

Five mandatory elements (ALL must hold):

1. **Trigger cadence (5 trigger types).** Pre-build (blocking, <24 h freshness), pre-flash (blocking, <24 h), pre-tag (blocking, <6 h), daily 09:00 UTC default (non-blocking sweep), post-deployment burn-in T+24 h (non-blocking).
2. **Three-source query.** Crashlytics + Analytics + Performance — ALL three. Skipping one source is a §11.4 PASS-bluff (Firebase reports a regression on the skipped axis and we never see it).
3. **Issues.md output.** Every "problematic" finding MUST map to either (a) a new Issues.md entry, OR (b) a recognised existing entry via the dedup algorithm. Both paths produce a full Firebase Console URL in the entry's metadata so any future agent can re-fetch the raw data without rebuilding context.
4. **Three-tier deduplication.** Tier 1 (exact Firebase Issue-ID match) → Tier 2 (stacktrace-similarity cluster hash: top-3 non-generic frames normalised + crash class → SHA-256 first 8 hex chars) → Tier 3 (monthly operator merge review). Multiple Issues.md entries pointing at the same underlying root cause is a §11.4 violation (operator chases ghosts; root cause stays unfixed).
5. **Comprehensive root-cause analysis.** Every Firebase-sourced Issue carries the §11.4.4(a) systematic-debugging output — Phase 1 evidence, Phase 2 pattern, Phase 3 root-cause hypothesis (UNCONFIRMED until §11.4.43 RED test reproduces), Phase 4 fix direction. Comment-only entries ("crash in App X") are PASS-bluff stubs and FAIL the Issue-xref gate.

Severity classification: Crashlytics FATAL impactedUsers > 10 last 7d → Critical; 1–10 → Major; 0 but historic → Low. NON-FATAL

100/day → Major; 10–100 → Low; < 10 → Informational. ANR 1% sessions → Major; 0.1–1% → Low. Performance KPI regression 20% → Major; 10–20% → Low. Analytics funnel-drop > 50% → Major; 25–50% → Low; feature_used regression > 75% → Major. Composes with §11.4.16 Type assignment.

Gates:

- CM-COVENANT-114-47-PROPAGATION — anchor present across every CLAUDE.md / AGENTS.md in the covenant file set.
- CM-AF-FIREBASE-REVIEW-CADENCE —
scripts/firebase/review_round.sh exists, executable, contains the 7-stage pipeline literals + dedup algorithm + severity table.
- CM-AF-FIREBASE-ISSUE-XREF — every Issues.md entry whose
Source: is Firebase Crashlytics / Firebase Analytics /
Firebase Performance carries **Firebase Issue IDs:** + **Firebase URL:** + (**Stacktrace Cluster Hash:** for Crashlytics OR **KPI:** for Performance OR **Funnel:** for Analytics).

Paired mutations strip the anchor / rename the review_round function / remove the Firebase-URL field from the template → respective gate FAILs.

Composition. §11.4.4 (STOP-on-discovery), §11.4.4(a) (systematic-debugging), §11.4.6 (no-guessing — Firebase data IS captured evidence), §11.4.7 (demotion-evidence — a previously- Fixed Issue resurfacing in Firebase IS positive captured evidence the fix didn't hold; reopen attribution = "AI: captured-evidence-contradicts" per §11.4.34), §11.4.10 (credentials — never log the bearer token), §11.4.12 (Issues_Summary sync), §11.4.14 (cleanup of /tmp/firebase_review_* work-dirs), §11.4.15 (Status: Queued), §11.4.16 (Type per severity), §11.4.34 (Reopened-Details on Firebase-resurface), §11.4.42 (implicit step 2 inclusion), §11.4.43 (RED test per stacktrace before fix lands), §11.4.44 (revision header on every Firebase-sourced Issue), §11.4.45 (Firebase Status.md at docs/firebase/status/Status.md), §11.4.46 (validation pass consults latest Firebase delta).

No escape hatch. No `--skip-firebase-review` ,
`--no-issue-from- firebase` , `--firebase-review-not-applicable` flag is permitted. The operator MAY filter by minimum severity (`--severity-min major`) to reduce noise but the pass itself MUST execute.

Classification: universal (per §11.4.17) — every consuming project that ships to real end-user devices benefits from the same review cadence and the same dedup algorithm.
11 4 48
2026 05 18

§ . . – UI-Driven Video Testing Mandate (User mandate, – –)

Forensic anchor — verbatim user mandate (2026-05-18):

"We MUST make sure that all video playback testing with 2nd display is performed by fully automatically using the application's UI / UX, with full navigation, choice of video and clicks to real UI buttons to play or switch / stop videos being played! Maybe direct execution of Intents or Broadcasts does not have reported issues! We MUST test ALL WAYS users or Systems may / could play some video contents or streams! For all applications, all video and audio stream types! All supported codecs! EVERYTHING verified with 2nd display on D3 device and every other device with connected 2nd display and Avrus for proper codec(s) used in its Web Interface / Dashboard! Write as much as possible new tests for all supported test types!"

Why the existing video-test set is insufficient. The legacy

`test_video_routing.sh` (v33), `test_video_secondary_display.sh` ,
`test_video_playback_automation.sh` (v1) plus the 53 `per_app_video` wrappers all dispatch playback via `am start -a android.intent.action.VIEW` or via `MediaSession cmd media_session play` — surfaces that skip the app's own content-selection UI, player-surface allocation, `MediaSession` lifecycle (start/pause/stop), and back-button cleanup. A defect like §CL routing-race or §CM frozen-frame mid-back-press cannot reproduce via Intent dispatch because Intent dispatch re-launches the app cold each time and `am force-stop` is a clean cleanup path that masks the leak.

Operative rule. Every video-playback test that asserts secondary-display routing MUST traverse the user-equivalent UI path. Mandatory 5-surface coverage per video app:

1. **Surface A — Launch via launcher icon.** `am start -n <launcher-pkg>/<launcher-activity>` then `uiautomator dump` → locate app icon tile → `input tap X Y`. NOT `am start <video-pkg>/<MainActivity>` (skips launcher).
2. **Surface B — Navigate to content list.** Per-app driver script knows the app's home-tab → content-tab swipe / tap sequence. Asserts at least one content tile is visible via `uiautomator dump` content-desc / text match.
3. **Surface C — Select + play.** Tap a specific content tile (driver records the picked tile's content-desc as captured evidence). Wait for player surface visible.
4. **Surface D — In-app control interaction.** Pause via in-app pause button (not `input keyevent KEYCODE_MEDIA_PAUSE`), resume via in-app play button, in-app rapid switch to a second video (covers §CL).
5. **Surface E — Stop via back button.** `input keyevent KEYCODE_BACK` to exit player. Assert §11.4.14 cleanup: `VOM activeDecoder == null`, secondary surface cleared, no orphan `MediaSession`.

App-coverage matrix: every video-capable app in `PRODUCT_PACKAGES` MUST have a UI-driven driver script set (Layer 2 per Section 7 design) AND at least one scenario that uses it. Today: 15 apps covered fully in the initial batch (VK Video, MPV, Lampa as fully-implemented; remaining 12 as templated stubs per Section 8). Expansion to all 53 video-app entries is queued as follow-up work under §O Issues.md per-app matrix.

Stream-type matrix: progressive HTTP / HLS / DASH / RTMP / file- local / DRM-protected. Per app, the driver script tags which stream types its picked content exercises. Aggregator emits a coverage matrix per cycle so gaps are visible.

Codec matrix: H.264 / H.265 / VP9 / AV1 / MPEG-2 / MP4V (video) + AC-3 / E-AC-3 / TrueHD / DTS / DTS-HD / MLP / Opus / AAC (audio). Per scenario, the codec assertion is made via `dumpsys media .metrics` (decoder name in event log) AND via Arvus codec-state probe (audio side) per §11.4.13.

Secondary-display verification. On any device with HDMI-A-1 attached (detected via `dumpsys display | grep "Display 2"`):

- Captured-evidence: `dual_display_record.sh` ON for both displays per §11.4.5
- Assertion: `ffprobe -count_frames` on secondary mp4 reports

N frames within 3 s of play-action timestamp

- Anti-assertion: `ffprobe -count_frames` on primary mp4 reports ONLY launcher/home pixels during playback window (no video-decoder output)
- VOM cross-check: `dumpsys video_output_manager` shows `activeDecoder != null` during playback and `== null` after stop within 3 s

When secondary absent (D2 default topology): SKIP per §11.4.3 with explicit reason "topology: no secondary display attached". NEVER FAIL on missing topology.

Arvus codec-state mandate. For every test asserting an audio codec (AC-3, E-AC-3, TrueHD, DTS, etc.) — composed with §11.4.13

- §CG:
- During playback window, call `arvus_probe_codec_state` ≥ 3 times at 1 s intervals (transients fade by sample 2-3)
- Call `arvus_screenshot_capture` per §CG to attach visual evidence of the dashboard
- Assert reported codec matches expected via `arvus_assert_codec_format <regex>`
- ARVUS_HOST unreachable (ABK4 not joined / sink off) → SKIP per §11.4.3, never FAIL

Pre-build gates:

- `CM-COVENANT-114-48-PROPAGATION` — asserts anchor literal present in every CLAUDE.md / AGENTS.md across parent + 10 owned submodules + HelixQA dependencies (42-file scan).

- `CM-AF-UI-DRIVEN-VIDEO-COVERAGE` — asserts the directory `device/rockchip/rk3588/tests/ui_driven/` exists, contains `lib/ui_driver.sh` reference + scenarios/ subdir + per-app driver subdirs for at least the 3 reference apps (vk_video, mpv, lampa), and that `scripts/testing/run_ui_driven_video_suite.sh` is executable.
- Paired mutations: deleting `ui_driver.sh` → CM-AF-UI-DRIVEN- VIDEO-COVERAGE FAILs; stripping the anchor literal → CM-COVENANT-114-48-PROPAGATION FAILs.

No escape hatch. No `--use-intent-shortcut / --skip-ui-traverse / --legacy-intent-mode` flag exists. The discipline exists because Intent-based dispatch is exactly the class of test that PASSes while users hit real defects (§11.4 forensic anchor). Operators who feel UI-driven tests are "too slow" should land scenario-level parallelism (multi-device fan-out) rather than bypass the discipline.

Composition.

- §11.4.3 — topology dispatch (secondary present vs absent)
- §11.4.5 — captured-evidence quality (every UI tap captured)
- §11.4.13 — Arvus codec-state mandate
- §11.4.14 — playback cleanup verified via UI back-button path
- §11.4.43 — RED-test-first discipline (initial UI tests RED against §CL/§CM)
- §11.4.44 — revision header on every plan/driver doc
- §CG — Arvus dashboard screenshot

Classification: ^{11 4 49}universal (per §11.4.17). Applies to every project that consumes the ^{2026 05 18}constitution submodule AND ships any video-capable Android app.

§ . . — Dual-Approach Testing Mandate
(User mandate, — —)

Forensic anchor — verbatim user mandate (2026-05-18):

"Kinopoisk playback of 5.1 audio supported movies MUST BE done via UI / UX full automation and via execution of Intents (or Broadcasts) directly both! We could have shared tests base and specialized part which will run / play the movie with properly chose audio track (5.1). Document everything up to the smallest details! Create as much as needed new tests for all supported test types and Challenges!"

Composition with §11.4.48. §11.4.48 mandated UI-driven traversal for every video routing test, eliminating the Intent-only PASS-bluffs. §11.4.49 REFINES rather than replaces: every feature test ships in BOTH variants — UI-driven AND Intent-driven — over a shared assertion base. Either alone is a §11.4 PASS-bluff for the OPPOSITE half of the stack. UI catches app-side bugs (content selection, player surface allocation, MediaSession lifecycle, in-app overlays, back-button cleanup). Intent catches framework/system-server bugs (Intent extras parsing, MediaCodec.configure hook MIME routing, IVideoOutputManager bind timing, broadcast permission gating, headless cron automation paths).

Operative rule — 5 mandatory elements:

1. **Both variants required.** Every feature test exercising a user-visible behaviour MUST ship both `<feature>_ui.sh` (uiautomator-based, §11.4.48 surfaces A–E) AND `test_<feature>_intent.sh` (`am start --es / am broadcast` -based). Either alone is forbidden.
2. **Shared assertion base.** Codec-state assertions, captured- evidence collection (screen-recording, ffprobe, RMS amplitude analysis), Arvus codec-state probe + dashboard screenshot, and §11.4.14 cleanup MUST be implemented in a single POSIX-sh library (`tests/lib/dual_approach_test_base.sh`) and reused by BOTH variants of every dual-approach test. Code duplication in the shared layer is forbidden.
3. **Specialised driver code.** UI variant uses §11.4.48 per-app driver scripts under `tests/ui_driven/<app>/` . Intent variant uses `am start --es content_id <id> --es audio_track_id <track> / am broadcast -a <action>` directly. Each variant's specialised layer is responsible ONLY for "how the playback gets started"; everything downstream of "playback is now happening" routes through the shared base.

4. **Comprehensive documentation per test.** Every dual-approach test set ships with: (a) §11.4.44 revision header in BOTH variant scripts, (b) per-feature contract under `docs/dual_approach/ <feature>.md` describing the captured-evidence pairing and operator-side verification, (c) `Issues.md` / `Fixed.md` entry cross-linking both variant paths.
5. **Kinopoisk 5.1 EAC3 is the canonical first implementation.** The shared base + Kinopoisk 5.1 UI variant + Kinopoisk 5.1 Intent variant land in the same batch as this mandate. Subsequent features (Netflix Dolby Atmos, MPV DTS-HD, Lampa+TorrServe HEVC, VLC FLAC) port to the same pattern. The §CN subagent's fix to the underlying decoder pipeline is the FIRST GREEN target of these tests; both variants are RED per §11.4.43 until §CN lands.

Captured-evidence directory contract. Both variants write to mirror-structured directories `qa-results/dual_approach/<F>/<run-ts>/{ui,intent}/`. Identical filenames; orchestrator diffs the two `result.json` files. Status mismatch (UI=PASS but Intent=FAIL or vice-versa) is itself a finding — it pinpoints which half of the stack contains the bug.

Pre-build gates:

- `CM-COVENANT-114-49-PROPAGATION` — anchor literal in every `CLAUDE.md` / `AGENTS.md` across parent + 10 owned submodules + HelixQA dependencies (42-file scan).
- `CM-AF-DUAL-APPROACH-COVERAGE` — `tests/lib/dual_approach_test_base.sh` exists + sources `anti_bluff` / `ui_driver` / `credentials` + exports `dat_init` / `dat_start_capture` / `dat_assert_codec_state` / `dat_assert_video_frames` / `dat_assert_audio_channels` / `dat_arvus_dashboard_capture` / `dat_cleanup` / `dat_report_finding`.
- `CM-AF-KINOPOISK-5-1-DUAL-COVERAGE` — both `tests/ui_driven/kinopoisk/kinopoisk_5_1_play_movie_ui.sh` AND `tests/test_kinopoisk_5_1_play_movie_intent.sh` exist + both source the shared base + both reference EAC3 codec + both assert 5.1 / 6-channel.
- Paired mutations (3): strip anchor → propagation gate FAILs; delete UI variant → coverage gate FAILs; delete Intent variant → coverage gate FAILs.

No escape hatch. No `--ui-only` / `--intent-only` / `--skip-dual` flag exists. The discipline exists because either-alone is the PASS-bluff pattern §11.4 specifically prohibits. Operators who feel dual variants are "too slow" should land scenario-level parallelism rather than bypass the discipline.

Composition.

- §11.4.3 — topology dispatch (Arvus reachable, secondary present)
- §11.4.5 — captured-evidence content-quality analysis
- §11.4.6 — UNCONFIRMED tagging for un-verified Kinopoisk element IDs
- §11.4.10 — credentials never echoed, per-device per-service loader
- §11.4.13 — Arvus codec-state mandate (out-of-band evidence)
- §11.4.14 — playback cleanup via EXIT trap in shared base
- §11.4.17 — universal classification
- §11.4.43 — RED-first TDD (both variants RED until §CN fix lands)
- §11.4.44 — revision header on every variant script + design doc
- §11.4.48 — UI-driven traversal (this MANDATE refines, not replaces)
- §CG — Arvus dashboard screenshot via Presenter receiver
- §CB — credentials loader

Classification: universal (per §11.4.17). Applies to every project that consumes the constitution submodule AND ships any Android app whose user-visible behaviour can be triggered through both UI and Intent/Broadcast paths.

§ 11.4.50 — Deterministic Consistency Mandate (User mandate, 2026-05-18)

Forensic anchor — verbatim user mandate (2026-05-18):

"Our anti-bluff / bluff-proofed / proof-driven work MUST satisfy the following: not a single feature, System or application flow, use case, edge case or procedural action(s) MUST NOT partially work, or sometimes work and sometimes not! There is only one truth that MUST BE fulfilled - no matter how many times we repeat all these scenarios with variations in data, System state(s) and speed of execution, results MUST BE consistent and successful without exception! No false positives or bluffing of any kind is allowed! We MUST add as much full automation tests which will use real System and Applications UI and UX with all flows, use cases and edge cases as needed or as much as it is possible!"

Why this anchor exists. §11.4.7 already forbids the vocabulary `intermittent` / `transient` / `flake` in closure narratives, but that's a textual ban — operators could still let a test "PASS once, FAIL once, PASS again" and report only the first PASS. §11.4.50 closes that gap MECHANICALLY by requiring every PASS to come from N identical iterations rather than from a single observation.

Operative rule — 5 mandatory elements:

1. **N-iteration deterministic outcome.** Every test that PASSes MUST have been executed N times (default N=3, N=10 for cycle- validation suites) against the same firmware MD5 + same device
 - same topology and produced IDENTICAL PASS in every iteration. A test whose outcome diverges across iterations is auto-FAIL per this mandate — there is no "first PASSEd therefore X was a flake" path. §11.4.7's expanded forbidden-vocabulary list is enforced mechanically here.
2. **Edge-case + flow + use-case coverage.** Every public API path (Activity / Service / Broadcast receiver / ContentProvider URI / IPC interface / JNI entry / sysprop write / sysfs node / init.rc trigger) MUST have ≥ 1 dedicated test that drives it. Untested paths surface in a feature-coverage-matrix audit and block release at the 99% threshold (ratchet sequence 70 $\rightarrow$ 85 $\rightarrow$ 95 $\rightarrow$ 99 mirrors §11.4.18 the existing project-side docs-coverage ratchet).
3. **Reliability check helper.** Project anti-bluff helper library ships

```
ab_run_n_times <test_name> <N> <fn> [args...] :
```

 - Loops N times, captures exit code + evidence-hash per iter

- Asserts all N exit codes identical AND all N evidence-hashes identical
- On any divergence: `ab_fail` with full N-iteration report
- Per §11.4.7 forbidden vocabulary: NO operator-facing escape path converts a divergent N-iter run into PASS

4. **Feature-coverage-matrix audit.** Project ships `scripts/testing/feature_coverage_audit.sh` :

- Walks source layers (every public API surface)
- Walks test layers (`test_*.sh` references)
- Emits coverage report at `qa-results/feature_coverage/ <ISO-UTC>/coverage.tsv`
- Pre-build gate enforces minimum threshold; threshold ratchets up each phase (70 → 85 → 95 → 99)

5. **Composes with every prior anti-bluff anchor.** §11.4.50 does NOT replace §11.4.43 / §11.4.48 / §11.4.49 — it adds the N-iter consistency dimension across all of them. RED-first TDD (§11.4.43) still applies: every new test starts RED and becomes GREEN after the fix lands. UI-driven (§11.4.48) and dual-approach (§11.4.49) still apply: every test ships in both variants AND each variant passes N iterations identically.

Pre-build gates:

- `CM-COVENANT-114-50-PROPAGATION` — anchor literal `§11.4.50` present in every CLAUDE.md / AGENTS.md across parent + 10 owned submodules + HelixQA dependencies (≥ 42 -file scan; phase 1 uses 5-canonical-file scan matching §11.4.49 propagation pattern).
- `CM-AF-RELIABILITY-CHECK-WIRED` — `ab_run_n_times` function literal present in `device/rockchip/rk3588/tests/lib/anti_bluff.sh` AND at least 3 on-device tests source-reference it. The 3-test floor ratchets upward each phase.
- `CM-AF-FEATURE-COVERAGE-MATRIX` — `scripts/testing/feature_coverage_audit.sh` exists + executable + has run within the last 7 days + most-recent coverage report meets the current threshold.

Paired mutations (3):

- Strip the §11.4.50 anchor literal from constitution/CLAUDE.md via `sed -i 's|11\.4\.50|11.4.MUTATED|g'` → CM-COVENANT-114-50-PROPAGATION MUST FAIL.
- Rename `ab_run_n_times` to `ab_run_n_times_DISABLED` in `anti_bluff.sh` via targeted `sed` → CM-AF-RELIABILITY-CHECK- WIRED MUST FAIL.
- Move `scripts/testing/feature_coverage_audit.sh` aside via `mv` → CM-AF-FEATURE-COVERAGE-MATRIX MUST FAIL.

No escape hatch. No `--allow-flake`, `--first-pass-suffices`, `--skip-n-iter`, `--skip-coverage-audit` flag exists. The discipline exists because the user mandate is unambiguous: "results MUST BE consistent and successful without exception."

Composition.

- §11.4.1 — FAIL-bluffs forbidden
- §11.4.2 — captured-evidence (every iteration captures its own)
- §11.4.5 — quality analysis (per-iteration content compared)
- §11.4.6 — no guessing (UNCONFIRMED required for any iteration- budget assumption)
- §11.4.7 — forbidden flake/intermittent vocabulary (enforced mechanically by this mandate)
- §11.4.43 — RED-first TDD (every new test starts RED, N-iter applies across the GREEN transition)
- §11.4.46 — validate-before-suite (N-iter applied at single-test scope first)
- §11.4.48 — UI-driven traversal (each UI traversal MUST be N-iter-consistent)
- §11.4.49 — dual-approach (both variants of every dual-approach test MUST pass N iterations identically)

Classification: universal (per §11.4.17). Applies to every project that consumes the constitution submodule.

§ . . – Live-ADB-First Maximization Mandate (User mandate, – –)

Forensic anchor — verbatim user mandate (2026-05-18):

"Was it possible to max the fix and test it via ADB or it can be only done by making the fix and reflashing for validation and verification? If such option exists, we should always make the best of it! Making fix on live devices via ADB → Once done commit and push changes, rebuild System and reflash → Validate and verify everything (with the test we write as well). Make sure this idea(s) is (are) part of our root (constitution Submodule) Constitution, AGENTS.MD and CLAUDE.MD if they are not already! EVERY DETAIL IS IMPORTANT!!!! Not a single idea or idea's detail can be ignored, skipped, relativized or bluffed!"

§11.4.51 REFINES §11.4.43 step 2 ("LIVE-ADB-PROBE — try the fix on the running device first") with mechanical enforcement: a per-file- class decision matrix, a classifier helper, and a commit-message footer literal.

Operative rule (5 mandatory elements):

1. **Classify every fix** by rebuild-requirement using the project's per-file-class decision matrix. No guessing (§11.4.6). Every file in the diff cites the matrix row that classified it.
2. **If LIVE_ADB_TESTABLE**, the operator **MUST** apply the fix to the running device via the appropriate `adb` channel first (`adb push` for scripts, `setprop persist.*` for runtime properties, `mount -o remount,rw /vendor` for boot scripts, `pm install -r` for APKs built locally via gradle). Run the §11.4.43 RED test against the live-probed device, capture positive-evidence PASS, THEN commit source-side + rebuild + reflash as belt-and-suspenders re-validation. Commit message footer **MUST** state `LIVE_ADB_VALIDATED: yes` with one-line description of the live-probe sequence.
3. **If REQUIRES_REBUILD**, the operator proceeds directly to source-side fix + rebuild + flash. Commit message footer **MUST** state `REQUIRES_REBUILD: <reason>` citing the matrix row that classified the file. Categories of `REQUIRES_REBUILD`: kernel, AOSP framework Java / AIDL, native C++ inside

APEX, sepolicy, init.rc, `ro.*` build properties (immutable post-boot),
`Android.bp` / `Android.mk` / `BoardConfig.mk` , XML resource overlays,
codec XML inside APEX, ramdisk-bundled artifacts.

4. **If mixed batch**, commit footer states `LIVE_ADB_VALIDATED: partial` and enumerates per-file classification. Per §11.4.9 batching is preferred; live-probe the testable files first, then batch with the rebuild-required ones into a single rebuild.
5. **Helper script enforces classification mechanically.** The project provides `classify_fix_rebuild_requirement.sh` which walks `git diff --name-only` , looks up each file against the matrix, and emits the recommended commit-message footer. No operator-facing override flag converts an unclassified path to `LIVE_ADB_TESTABLE` by default — unmatched paths classify as `REQUIRES_REBUILD: unmatched-path` (safe default per §11.4.6).

Per-file-class decision matrix (canonical):

File class	Rebuild required?	Reason
<code>*test_*.sh</code> / <code>tests/lib/*.sh</code> (on-device)	NO	adb push to <code>/data/local/tmp/tests/</code>
Host-side test scripts	NO	host script — no device interaction
<code>atmosphere-*.sh</code> boot scripts	DEPENDS	adb push to <code>/vendor/bin/</code> after remount
Markdown docs / Constitution / plans / Issues / Fixed / CONTINUATION	NO	host docs — no firmware impact
Forked-player Kotlin (Presenter, MPV, SmartTube, TorrServe, Lampa, VLC, etc.)	YES via gradle local build + <code>adb install -r</code>	APK re-link required but no full firmware rebuild
Framework Java (<code>frameworks/base/**/*.java</code>)	YES	system_server requires reboot at minimum
AIDL (<code>*.aidl</code>)	YES	stub generation requires recompile
Native C++ (<code>external/**/*.cpp</code> , <code>frameworks/native/**/*.cpp</code>)	YES	builds into APEX or system library
APEX libraries (<code>vendor/**/lib/**/*.*.so</code> inside APEX path)	YES	squashfs read-only by Mainline integrity
Kernel sources (<code>kernel-5.10/**</code>)	YES	requires kernel rebuild + boot.img
sepolicy (<code>*.te</code>)	YES	requires policy build
<code>init.rc</code>	YES	parsed at boot only
<code>ro.*</code> build properties	YES	immutable post-boot
<code>persist.*</code> runtime properties	NO	<code>setprop persist.*</code> mutable

File class	Rebuild required?	Reason
XML resource overlays	YES	RRO recompile
<code>media_codecs_*.xml</code>	YES	APEX-bundled
<code>display_settings.xml</code> (user-data)	DEPENDS	adb push possible if SELinux permits
<code>Android.bp</code> / <code>Android.mk</code> / <code>device.mk</code> / <code>BoardConfig.mk</code>	YES	rebuilt-into-image artifacts
Test fixture binary assets	NO	adb push to <code>/data/local/tmp/</code>

Pre-build gates: `CM-COVENANT-114-51-PROPAGATION` (anchor across canonical files) + `CM-AF-CLASSIFY-FIX-HELPER-EXISTS` (helper present + sentinel literals) + `CM-AF-LIVE-ADB-FIRST-COMMIT-MARKER` (advisory WARN that scans recent commits for the footer literal). Three paired meta-test mutations.

Composition.

- §11.4.43 — TDD-fix workflow (§11.4.51 REFINES step 2 with the mechanical classifier).
- §11.4.9 — batch-source-fixes-before-rebuild (compose: live-ADB- test individually, then batch-commit).
- §11.4.6 — no guessing (each classification justified by matrix row; unmatched paths default to rebuild-required, never silent testable).
- §11.4.46 — validate-recent-work-before-post-flash (live-probe IS implicit pre-validation).
- §11.4.48 — UI-driven tests (live-probe CAN use uiautomator over adb shell).
- §11.4.49 — dual-approach (live-probe applies to both UI and Intent variants).
- §11.4.50 — deterministic consistency (live-probe runs N iterations on the live device before commit).

No escape hatch — no `--skip-classify` , `--assume-rebuild` , `--no-footer-required` flag exists. The discipline exists because the user mandate is unambiguous: "EVERY DETAIL IS IMPORTANT!!!! Not a single idea or idea's detail can be ignored, skipped, relativized or bluffed!"

Classification: universal (per §11.4.17). Applies to every project that consumes the constitution submodule — the matrix itself MAY be project-specific (each consumer adapts the file-class list to its own source tree), but the mandate to classify-before-commit + the LIVE_ADB_VALIDATED / REQUIRES_REBUILD footer literals + the classify-helper-script + the three pre-build gates are universal.

2026 05 18

§ . . — Autonomous-Validation Mandate (User mandate, — —)

Forensic anchor — verbatim user mandate (2026-05-18):

"Make sure we have full automation tests which will do all this work in full automation! IMPORTANT: Make sure that all existing tests and Challenges do work in anti-bluff manner — they MUST confirm that all tested codebase really works as expected! We had been in position that all tests do execute with success and all Challenges as well, but in reality the most of the features does not work and can't be used! This MUST NOT be the case and execution of tests and Challenges MUST guarantee the quality, the completion and full usability by end users of the product! This MUST BE part of Constitution of our project, its CLAUDE.MD and AGENTS.MD if it is not there already, and to be applied to all Submodules's Constitution, CLAUDE.MD and AGENTS.MD as well."

Why this anchor exists. §11.4.25 (full-automation coverage) and §11.4.27 (no-fakes-beyond-unit-tests + 100%-test-type-coverage) and §11.4.39 (per-feature on-device end-user validation) collectively mandated that every user-facing feature has automation tests with captured runtime evidence. They did NOT explicitly forbid operator-attended-only validation paths — i.e., a feature whose ONLY proof of working state requires a human to physically present at the device, drive UI manually, and observe outcomes. Phase 39.§CN exposed the gap: §CN's Kinopoisk 5.1 EAC3 closure was structurally fixed (libavutil bundled in APEX, six decoders register) yet end-user-validation degraded to "operator drives remote control while test polls" because Kinopoisk's TV-Compose UI is not accessibility-instrumented and uiautomator returns near-empty hierarchy. That validation path is

operator-attended, non-reproducible in CI, and cannot prove deterministic consistency (§11.4.50) across N iterations. The User mandate elevates "full automation" from "preferred" to "mandatory": every user-facing PASS MUST have at least one captured-evidence path that does NOT require operator presence.

Operative rule (4 mandatory elements):

1. **Every user-facing feature MUST have at least one autonomous validation path.** "Autonomous" means: the validation runs end-to-end via `adb shell` + scripted automation, produces captured runtime evidence per §11.4.5, and reaches a PASS / FAIL verdict WITHOUT a human present to drive UI, observe screen, or make decisions. Operator-attended tests are SUPPLEMENTARY, never PRIMARY. A feature whose ONLY validation path is operator-attended is a §11.4.52 violation regardless of how carefully the operator's captures match §11.4.5 evidence requirements — the path does not scale to CI, does not run on every commit, does not survive operator unavailability, and produces the exact "tests pass but feature doesn't work for users" failure mode §11.4 specifically forbids.
2. **Acceptable autonomous-validation paths** (any one or more — non-exhaustive):
 - **Programmatic instrumentation APK** — a small platform-signed app that exercises the under-test surface directly via SDK APIs (e.g., `MediaCodec.createDecoderByName`, `MediaCodec.configure`, `AudioTrack` writes) and writes a structured JSON result file to a discoverable path. Closes UI-instrumentation gaps where the target app is not accessibility-instrumented.
 - **Headless intent dispatch + state poll** — `am start --es /`
`am broadcast --es` to dispatch the under-test code path, then poll `dumpsys`, `/proc/<pid>/maps`, `media.metrics`, `/sys/...`, kernel sysfs, or the integration's network probe for the expected state transition.
 - **ADB-driven uiautomator** — applicable ONLY if the target app exposes accessibility nodes the dump can resolve. MUST verify via `uiautomator dump | grep -c clickable=true ≥ 1` before claiming uiautomator coverage is real. A near-empty hierarchy is itself the captured-evidence proving UI-driven automation is INFEASIBLE for that target and the test MUST fall back to instrumentation-APK or headless-intent paths.

- **Network-side sink probe** — for features that surface on a network-accessible peripheral (e.g., Arvus HDMI dashboard at `http://<sink>/` , Sonos REST API, AirPlay receiver), the sink's own report counts as autonomous out-of-band evidence per §11.4.13.
- **HelixQA autonomous QA session** — for projects that have HelixQA fully incorporated per §11.4.27, the orchestrator's end-to-end session output counts as autonomous evidence.

3. **Per-feature classification + tracking.** Every entry in the project's coverage ledger (§11.4.25) MUST classify each feature's autonomous-validation path as one of: `AUTONOMOUS_VERIFIED` (path exists + last-run-green per §11.4.46), `AUTONOMOUS_DESIGNED` (path exists but RED per §11.4.43), `OPERATOR_ATTENDED_ONLY` (path requires human presence — release blocker until promoted to one of the other states), or `NOT_APPLICABLE` (e.g., a CLI tool that has no UI surface and the unit + integration tests already constitute autonomous proof). `OPERATOR_ATTENDED_ONLY` rows MUST cite a tracked work item per §11.4.15 + §11.4.16 that designs the autonomous-path migration. Closing the work item requires the migration to land, not just the operator capture.

4. **Anti-bluff for autonomous paths themselves.** An autonomous path that returns PASS without exercising the user-visible behaviour is a §11.4.52 violation just as severe as the original §11.4 PASS-bluff pattern. The autonomous path MUST produce positive captured evidence per §11.4.5 (audio: channel count + sample rate + glitch census; video: frame count + routing target

- frame health + obstruction census). A grep on a metadata field that doesn't prove behaviour is not an autonomous validation — it is a metadata bluff dressed up as automation. Paired meta-test mutations (§1.1) MUST catch the autonomous path's degradation to metadata-only or grep-only.

Composition.

- §11.4.25 — full-automation-coverage mandate. §11.4.52 is the strict expansion of invariant 1 + 2: "anti-bluff posture" and "proof of working capability" cannot be carried by operator-attended-only evidence — they MUST be carried by at least one autonomous path.

- §11.4.27 — no-fakes-beyond-unit-tests + 100%-test-type-coverage. §11.4.52 is the operational layer that closes the gap between "the project has full-automation tests" and "every PASS has an autonomous evidence path".
- §11.4.39 — per-feature on-device end-user validation. §11.4.52 refines: the on-device scenario MUST itself be autonomously drivable.
- §11.4.43 — TDD-fix-discipline. The autonomous path MUST exist in the RED state BEFORE the fix lands, then go GREEN after. An autonomous path authored after the fix is a §11.4 PASS-bluff.
- §11.4.48 — UI-driven video testing. §11.4.52 establishes that when uiautomator returns a near-empty hierarchy on the target app, UI-driven IS infeasible and the test MUST fall back to instrumentation-APK or headless-intent. The fallback is REQUIRED — not optional.
- §11.4.49 — dual-approach testing. The Intent variant in a dual-approach pair often IS the autonomous path when the UI variant requires operator assistance.
- §11.4.50 — deterministic consistency. Autonomous paths run N iterations naturally; operator-attended paths cannot satisfy the N-iteration discipline at scale.
- §11.4.51 — live-ADB-first maximization. The instrumentation-APK path is itself LIVE_ADB_TESTABLE (gradle build + `adb install -r`
 - run + read JSON result), so live-probe IS the design loop for the autonomous path.

Pre-build gates:

- CM-COVENANT-114-52-PROPAGATION — anchor literal §11.4.52 present in all 5 canonical files (`constitution/Constitution.md` , `constitution/CLAUDE.md` , `constitution/AGENTS.md` , plus the parent consumer's `CLAUDE.md` + `AGENTS.md`) PLUS the full per-consumer propagation (parent project + every owned submodule
 - nested submodules + HelixQA dependencies — same scan pattern as CM-COVENANT-114-51-PROPAGATION).
- CM-AF-AUTONOMOUS-PATH-PER-FEATURE — every entry in the consuming project's coverage ledger (§11.4.25) has a non-empty autonomous-classification column with value in {AUTONOMOUS_VERIFIED, AUTONOMOUS_DESIGNED,

OPERATOR_ATTENDED_ONLY, NOT_APPLICABLE} . OPERATOR_ATTENDED_ONLY rows MUST cite a tracked migration work item. (Project-side gate — consuming projects implement the ledger surface in their own pre-build test suite.)

Paired mutations (per §1.1):

- Strip §11.4.52 literal from constitution/Constitution.md → CM-COVENANT-114-52-PROPAGATION FAILS.
- Inject a feature row with OPERATOR_ATTENDED_ONLY and NO tracked migration work item → CM-AF-AUTONOMOUS-PATH-PER-FEATURE FAILS.

No escape hatch. No --allow-operator-attended-only , --skip-autonomous-path , --manual-validation-suffices flag exists. The discipline exists because the user mandate is unambiguous: "execution of tests and Challenges MUST guarantee the quality, the completion and full usability by end users". Operator-attended-only validation cannot guarantee that property at the release-gate layer — only autonomous paths can.

Classification: universal (per §11.4.17). Applies to every project consuming the constitution submodule. Each consumer adapts the autonomous-validation path catalogue to its own technology surface (Android: instrumentation APK + uiautomator + headless intent; Web: Playwright headless + API probe; CLI: subprocess

- stdout/stderr assertion; mobile native: device farm headless runner) but the mandate to provide at least one autonomous path per user-facing feature is universal.

2026 05 18

§ . . — Fixed_Summary parity mandate
(User mandate, — —)

Forensic anchor — verbatim user mandate (2026-05-18T17:55Z):

"Note: Just like for Issues we have Issues_Summary, for Fixed we MUST HAVE Fixed_Summary - like all other docs: ALWAYS in sync and up to date and ALWAYS exported into the PDF and HTML! Add this mandatory rule / constraint into the root (constitution Submodule) Constitution, AGENTS.MD and CLAUDE.MD."

Why this anchor exists. §11.4.12 already established that `docs/Issues_Summary.md` is the canonical short-form summary of `docs/Issues.md` and MUST always be regenerated + re-exported (HTML

- PDF) whenever the source changes. §11.4.19 added the column- alignment requirement so `Fixed_Summary.md` mirrors `Issues_Summary.md` shape. §11.4.53 closes the propagation gap: the parity discipline that §11.4.12 imposes on Issues / Issues_Summary MUST apply symmetrically to Fixed / Fixed_Summary. The mechanical behaviour already exists in `scripts/testing/sync_issues_docs.sh` (lines 98- 105 regenerate Fixed_Summary in stage 1b alongside Issues_Summary in stage 1a), but the mandate was never explicit in the Constitution — and an unstated rule is a §11.4 PASS-bluff risk: future refactors of `sync_issues_docs.sh` could drop the Fixed_Summary half without violating any canonical authority. §11.4.53 makes the existing behaviour explicit canonical authority so the gate set can defend it.

Operative rule. `docs/Fixed_Summary.md` is the symmetric short- form summary of `docs/Fixed.md` . It MUST be regenerated whenever `Fixed.md` changes. Its HTML + PDF exports MUST travel with the markdown (identical mtimes within `sync_issues_docs.sh` granularity). Stale exports are §11.4.53 violations regardless of whether the underlying `.md` is correct — an operator (or future agent) reading the HTML or PDF gets a divergent view of which items are closed, and the §12.10 CONTINUATION resumption guarantee silently breaks. Same discipline as §11.4.12 Issues_Summary applied to Fixed.md.

Generator. `scripts/testing/generate_fixed_summary.sh` is the canonical generator. It MUST exist + be executable + emit a markdown table whose header columns include `Status` and `Type` (per §11.4.19 column-alignment) so the `Fixed_Summary.md` shape mirrors `Issues_Summary.md` exactly.

Auto-sync wrapper. `scripts/testing/sync_issues_docs.sh` is the single operator-facing entry point. It already regenerates BOTH `Issues_Summary` AND `Fixed_Summary` in one shot (stages 1a + 1b), then exports HTML + PDF (stage 2), then colorizes per §11.4.23 (stage 3), then re-renders the PDFs from the colorized HTML (stage 3b). MUST be invoked after any edit to `Fixed.md` (just as §11.4.12 requires after any edit to `Issues.md`). Manual invocation of just the Issues half (`--issues-only` -style flag) is FORBIDDEN — the wrapper has no such flag and §11.4.53 prohibits adding one.

Sort order. Same pattern as §11.4.12 — by closure date DESC (most-recent-Fixed first), with §-letter / Fix-# secondary sort. Documented at the top of the generated file.

HTML + PDF travel-together. Per §11.4.12 + §11.4.44, the three file types (`.md` , `.html` , `.pdf`) MUST always have identical mtimes within `sync_issues_docs.sh` granularity. Stale exports violate §11.4.53 regardless of whether the underlying `.md` is correct.

Composition:

- §11.4.12 — `Issues_Summary` parity is the sibling rule §11.4.53 symmetrizes. §11.4.12 and §11.4.53 are the canonical pair: edits to one source-of-truth (`Issues.md` or `Fixed.md`) trigger regeneration of BOTH summaries via the same wrapper.
- §11.4.19 — atomic Issues → Fixed migration triggers `Fixed_Summary` regeneration. When an item closes (status `Fixed` (→ `Fixed.md`) / `Implemented` (→ `Fixed.md`) / `Completed` (→ `Fixed.md`) per §11.4.33), the operator moves the entry from `Issues.md` to `Fixed.md` atomically and runs `sync_issues_docs.sh` — both summaries refresh.
- §11.4.23 — visual-cue & grouping colorizer post-processes both `Issues_Summary.html` AND `Fixed_Summary.html`. Stale `Fixed_Summary.html` would carry pre-§11.4.23 styling — itself a §11.4.53 violation.
- §11.4.33 — type-aware closure-status vocabulary. `Fixed_Summary` MUST respect the three terminal values (`Fixed` (→ `Fixed.md`) , `Implemented` (→ `Fixed.md`) , `Completed` (→ `Fixed.md`)) and emit them literally in the Status column.

- §11.4.44 — revision-header mandate applies to `Fixed_Summary.md` exactly as it applies to every other tracked Markdown doc.
- §12.10 — CONTINUATION.md resumption guarantee depends on the divergent-summary problem NOT existing; §11.4.53 is the explicit symmetric closure of that risk for the Fixed half.

Pre-build gates:

- `CM-FIXED-SUMMARY-SYNC` — canonical artifact-level gate (6 invariants): (1) `docs/Fixed_Summary.md` exists; (2) `docs/Fixed_Summary.html` exists with `mtime ≥ Fixed_Summary.md mtime`; (3) `docs/Fixed_Summary.pdf` exists with `mtime ≥ Fixed_Summary.md mtime`; (4) `Fixed_Summary.md mtime ≥ Fixed.md mtime`; (5) `scripts/testing/generate_fixed_summary.sh` exists + executable; (6) `scripts/testing/sync_issues_docs.sh` invokes the generator. Pre-build FAILs if any invariant is violated. Partially overlaps with `CM-FIXED-COLUMN-ALIGNMENT` (§11.4.19 Phase 39.AV gate) and `CM-DOCS-EXPORT-SYNC` (§11.4.15 4-doc + 6-export scan), but is explicitly the `Fixed_Summary` derived-doc parity invariant for §11.4.53 enforcement.
- `CM-COVENANT-114-53-PROPAGATION` — anchor literal `§11.4.53` present across all 5 canonical files (`constitution/Constitution.md` , `constitution/CLAUDE.md` , `constitution/AGENTS.md` , parent consumer's `CLAUDE.md` + `AGENTS.md`). Same propagation pattern as `CM-COVENANT-114-51-PROPAGATION` and `CM-COVENANT-114-52-PROPAGATION` . Consumer projects propagate further (owned submodules, nested submodules, HelixQA deps) at their own per-project gate granularity.

Paired mutations (per §1.1):

- Strip `§11.4.53` literal from `constitution/CLAUDE.md` → `CM-COVENANT-114-53-PROPAGATION` FAILs.
- Move `scripts/testing/generate_fixed_summary.sh` aside (so the invariant-(5) executable-check fails) → `CM-FIXED-SUMMARY-SYNC` FAILs.
- Touch `docs/Fixed_Summary.md` to a date older than `docs/Fixed.md` → invariant-(4) violated → `CM-FIXED-SUMMARY-SYNC` FAILs.

No escape hatch. No `--skip-fixed-summary-sync` , `--issues-only` , `--summary-not-applicable` flag exists. The discipline exists because the user mandate is unambiguous: "like all other docs: ALWAYS in sync and up to date and

ALWAYS exported into the PDF and HTML". A divergent Fixed_Summary breaks operator + AI-agent ability to discover what is fixed vs what is still open — and that breakage IS the §11.4 "tests pass but reality differs" pattern at the documentation layer.

Classification: universal (per §11.4.17). Applies to every project consuming the constitution submodule that maintains an `Issues.md` / `Fixed.md` split (the Issues-tracking lifecycle is already a §11.4.15 universal mandate, so this clause inherits the same universal scope). Projects that do not maintain a `Fixed.md` (e.g., single-binary throwaway prototypes) are NOT_APPLICABLE; all projects that DO maintain a `Fixed.md` MUST land both gates and mutations within one cycle of adopting this Constitution version.

2026 05 19

§ . . — ATM-NNN ticket identifier mandate (User mandate, — —)

Forensic anchor — verbatim user mandate (2026-05-19):

"Every workable item in `Issues.md` / `Issues_Summary.md` / `Fixed.md` / `Fixed_Summary.md` MUST carry a stable, unique, auto-incremental ATM-NNN ticket identifier. ATM- prefix, monotonic, never renumbered, append-only — once assigned to an item it stays bound to that item for the lifetime of the project across every reopen, migration, or rename."

Why this anchor exists. §11.4.15 + §11.4.16 + §11.4.19 + §11.4.33 established the lifecycle vocabulary for tracked items, but the items themselves were identified only by §-letter or Fix-# — both of which are **unstable across migrations**. §-letter is heading-local; when an item moves from `Issues.md` to `Fixed.md` the §-letter may be reused for a different item. Fix-# is only assigned on closure for items that touched source code; tasks / features / pure-doc items have no stable identifier at all. Cross-references between commits, test logs, captured-evidence artifacts, Reopens.md history (§11.4.55), README.md doc-link rows (§11.4.57), and

external systems (Firebase Crashlytics Issue-IDs, Atlassian-style tracker rows) all need an identifier that is **monotonic, stable, and project-wide unique**. §11.4.54 closes that gap.

Operative rule. Every workable item in `docs/Issues.md` AND `docs/Fixed.md` MUST carry a `[ATM-NNN]` ticket identifier in its heading, in the form `## §X.Y. [ATM-NNN] <title>` (or `### §X.Y.` for nested items). NNN is a positive integer, zero-padded to at least 3 digits (`ATM-001` , `ATM-042` , `ATM-127` , ...). Identifiers are allocated by a **canonical helper script** (`scripts/testing/assign_atm_ticket_ids.sh`) that maintains an append-only state file (`scripts/testing/.atm_ticket_state.json`). Once assigned to an item, an ATM-NNN MUST NEVER be:

1. **Renumbered** — the bind is permanent.
2. **Reused** for a different item even if the original item is force-deleted (the helper's state file detects collisions).
3. **Decrementd** — counter is monotonic-only.
4. **Skipped intentionally** — gaps are forbidden in the assignment sequence (gaps indicate state-file corruption and trigger gate FAIL).

State file format. `scripts/testing/.atm_ticket_state.json` is a JSON-lines file with one record per allocated ID:

```
{"atm_id": "ATM-001", "heading_hash": "<sha256 of normalized heading>",  
  "type": "Bug|Feature|Task", "current_location": "Issues|Fixed",  
  "current_status": "<lifecycle value>", "reopens_count": <int>,  
  "created_at": "<ISO 8601 UTC>", "last_modified": "<ISO 8601 UTC>"}
```

The `heading_hash` is the canonical binding key — when the helper scans `Issues.md` / `Fixed.md`, it computes the normalized hash of each heading and looks up an existing ATM-NNN before allocating a new one. This protects against accidental renumbering on heading reflows / wording edits (the hash MUST stay stable; if the heading text changes substantially, the state file's `heading_hash` field is updated in place via an explicit migration command, NOT auto- rebound).

Summary-table column. `docs/Issues_Summary.md` and `docs/Fixed_Summary.md` MUST carry a leftmost `ATM ID` column showing the identifier. Generators (`generate_issues_summary.sh` , `generate_fixed_summary.sh`) MUST emit this column as the first data column so operators / agents can sort + filter on it.

Composition:

- §11.4.15 (Status), §11.4.16 (Type), §11.4.19 (column-alignment Issues ↔ Fixed), §11.4.33 (type-aware closure terminal values). All four pre-existing mandates compose with §11.4.54 — the ATM-NNN is the new universal join key.
- §11.4.12 + §11.4.53 (Issues_Summary + Fixed_Summary regen on source changes) — the helper MUST be invoked from `sync_issues_docs.sh` so newly added items receive ATM-NNN identifiers before summary regeneration.
- §11.4.55 (Reopens-history per-item docs) — the per-item Reopens.md lives at `docs/issues/<ATM-NNN>/Reopens.md` .
- §11.4.57 (README.md doc-link section) — every Issues / Fixed item may be linked by ATM-NNN in external references.
- §11.4.44 (revision header) — applies to the state file's audit log if one is maintained.

Pre-build gates:

- `CM-ATM-TICKET-IDS-COMPLETE` — every workable-item heading in Issues.md AND Fixed.md MUST carry an `[ATM-NNN]` token matching `\[ATM-[0-9]{3,}\]` . Headings without one FAIL the gate.
- `CM-ATM-TICKET-IDS-UNIQUE` — no two items share an ATM-NNN. The state file's record count MUST equal the union of Issues.md + Fixed.md `[ATM-NNN]` occurrences (no orphan IDs, no duplicates).
- `CM-ATM-TICKET-IDS-MONOTONIC` — the maximum ATM-NNN in the state file MUST equal `record_count` (no gaps).
- `CM-COVENANT-114-54-PROPAGATION` — anchor literal §11.4.54 present across canonical files (`constitution/Constitution.md` , `constitution/CLAUDE.md` , `constitution/AGENTS.md` , parent consumer's `CLAUDE.md` + `AGENTS.md`).

Paired mutations (per §1.1):

- Remove [ATM-NNN] from one Issues.md heading → CM-ATM-TICKET-IDS-COMplete FAILs.
- Duplicate an ATM-NNN in the state file → CM-ATM-TICKET-IDS-UNIQUE FAILs.
- Delete ATM-NNN=2 from a 5-record state file → CM-ATM-TICKET-IDS-MONOTONIC FAILs (gap at 2).
- Strip §11.4.54 literal from constitution/CLAUDE.md → CM-COVENANT-114-54-PROPAGATION FAILs.

No escape hatch. No `--skip-atm-assignment`, `--renumber`, or `--no-atm-id-required` flag exists. The discipline exists because unstable item identity is the exact failure mode that produces "the test was passing last week, what changed?" sessions where the operator cannot tell whether a given item is the same one tracked before or a coincidentally similarly-named successor.

Classification: universal (per §11.4.17). Applies to every project consuming the constitution submodule that maintains an Issues / Fixed split (cf. §11.4.53 scoping). Projects that do not maintain such a split are NOT_APPLICABLE.

2026 05 19

§ . . — Reopens-history tracking +
per-item Reopens.md doc (User mandate,
— —)

Forensic anchor — verbatim user mandate (2026-05-19):

"Add a Reopens-count column to Issues / Issues_Summary / Fixed / Fixed_Summary. For any item whose reopens-count > 0, create docs/issues/ATM-NNN/Reopens.md (+ HTML + PDF) with comprehensive reopen history + Fixed cycles. Each reopen MUST include By (AI / User), On (date), Reason, Evidence, and each Fixed-marking with its reasoning chain."

Why this anchor exists. §11.4.34 mandated that Reopened status carry a `**Reopened-Details:**` block — but that block lives in the **current** Issues.md heading and is overwritten on the next state change. A complete reopen / fix-cycle

audit trail across the lifetime of an item was not previously canonical. §11.4.55 makes the full history first-class: every item that has been reopened at least once gets a dedicated `Reopens.md` companion doc capturing the entire chronological timeline — every reopen, every closure, every reasoning chain — so an operator / agent investigating "why was this reopened five times" has a single, authoritative source to read.

Operative rule. Every workable item with `reopens_count > 0` MUST have a companion document at `docs/issues/ATM-NNN/Reopens.md` (plus its HTML + PDF exports per §11.4.44 + §11.4.53 sync discipline). The document MUST contain:

1. **Revision header** per §11.4.44 (`Revision` integer + `Last modified` ISO 8601 UTC).
2. **Item identification** — `ATM ID` , `Title` , `Type` , `Current Status` , `Current Location` (Issues.md / Fixed.md), link back to the live heading.
3. **Cycle counters** — `Total reopens` , `Total fixed cycles` (items can reopen multiple times, each followed by a re-closure).
4. **Chronological timeline** — one entry per state-change event, each entry capturing:
 - **By:** `AI` or `User` (per §11.4.34 source attribution).
 - **On:** ISO date.
 - **Event:** `Opened` | `Reopened` | `Fixed` | `Implemented` | `Completed` .
 - **Reason:** the closed-vocabulary value from §11.4.34 (`test-failed` / `manual-testing-detected` / `captured-evidence-contradicts` / `end-user-report` / `cycle-re-discovered` / `design-reconsidered`) for reopens; captured-evidence summary for closures.
 - **Evidence:** path to or short description of captured artefact (test log path, recording path, sink-probe report, operator quote).
 - **Outcome:** what the next state transition was.
5. **Reasoning chain for each closure** — for every `Fixed` / `Implemented` / `Completed` event, the reasoning chain that justified marking it closed (root-cause analysis link, captured-evidence under same-conditions per §11.4.7, gate / mutation pair that defends against regression).

6. **Most-recent state-change pointer** — explicit cross-reference to the current Issues.md / Fixed.md heading.

Summary-table column. Issues_Summary.md and Fixed_Summary.md MUST carry a `Reopens` column showing the integer count. When the count > 0 the cell MUST be a hyperlink to `docs/issues/<ATM-NNN>/Reopens.md`. Generators emit the column; the colorizer (§11.4.23) MAY apply a visual cue when `reopens > 2` (repeated-reopen signal — investigate before closing again).

Composition:

- §11.4.34 (Reopened-Details on current heading) is the **per-event** capture; §11.4.55 is the **per-item history** aggregation. Both layers MUST be in sync — the state-machine helper updates Reopens.md on every Reopened-Details parse.
- §11.4.54 (ATM-NNN) provides the stable identifier so the per- item doc has a permanent path even after Issues → Fixed migration.
- §11.4.44 (revision header) applies to every Reopens.md.
- §11.4.45 (Status.md per-integration) — the per-item Reopens.md is the item-scoped analogue of the integration-scoped Status.md.
- §11.4.53 (Fixed_Summary parity) — Reopens column appears on BOTH summary tables symmetrically.

Pre-build gates:

- `CM-REOPENS-DOC-EXISTS-WHEN-COUNT-GT-ZERO` — for every state-file record with `reopens_count > 0`, `docs/issues/<ATM-NNN>/Reopens.md` MUST exist (+ HTML + PDF). Missing doc FAILS.
- `CM-REOPENS-DOC-REVISION-HEADER` — every Reopens.md carries the §11.4.44 revision header.
- `CM-REOPENS-COL-IN-SUMMARIES` — both Issues_Summary.md and Fixed_Summary.md carry a `Reopens` column.
- `CM-COVENANT-114-55-PROPAGATION` — anchor literal §11.4.55 across canonical files.

Paired mutations (per §1.1):

- Delete `docs/issues/ATM-NNN/Reopens.md` for an item with `reopens_count=2` → `CM-REOPENS-DOC-EXISTS-WHEN-COUNT-GT-ZERO` FAILS.

- Strip revision header from a Reopens.md → CM-REOPENS-DOC-REVISION-HEADER FAILS.
- Remove Reopens column from Issues_Summary.md table header → CM-REOPENS-COL-IN-SUMMARIES FAILS.
- Strip §11.4.55 literal from constitution/CLAUDE.md → CM-COVENANT-114-55-PROPAGATION FAILS.

No escape hatch. No `--skip-reopens-doc` , `--collapse-history` , `--reopens-not-applicable` flag. The discipline exists because loss of reopen history is the exact failure mode that produces "why does this keep happening" sessions with no auditable answer.

Classification: universal (per §11.4.17). Applies wherever §11.4.34 applies (every project tracking Reopened status).

2026 05 19

§ . . – Status_Summary parity + two-audience format (User mandate, – –)

Forensic anchor — verbatim user mandate (2026-05-19):

"Every Status.md doc gets a Status_Summary parity companion ALWAYS in sync, exported to HTML + PDF. Two-page format: page 1 = non-developer audience (team-specific: audio team for audio Status, video team for video Status, etc.), page 2 = software engineer summary. Auto-generated after every main Status update."

Why this anchor exists. §11.4.45 established Status.md per integration domain — but Status.md targets engineers (cites gate names, §-letter references, commit hashes, captured-evidence file paths). Non-developer stakeholders (audio team, video team, QA leads, product reviewers) need a plain-language version of the same content, **always derived** from the engineering Status.md so the two views can never diverge. §11.4.56 introduces the symmetric companion doc with a two-audience layout.

Operative rule. For every `docs/<domain>/<integration>/ Status.md` that satisfies §11.4.45, a companion `docs/<domain>/<integration>/ Status_Summary.md` MUST exist with:

1. **Revision header** per §11.4.44.
2. **Page 1 — For the <team>** (audience-specific heading; the team name is derived from the integration domain — audio team for `docs/dolby/*` , video team for `docs/video/*` , etc.). Page 1 contains:
 - Plain-language summary of current state (1-3 paragraphs).
 - **What works** (1-3 bullets, end-user-visible behaviour).
 - **What's broken or pending** (1-3 bullets).
 - **Operator / team actions** if any.
 - NO code references, NO §-letter jargon, NO captured-evidence file paths, NO gate / mutation names.
3. **Page 2 — For software engineers** — same content density as Status.md but condensed:
 - §-letter references, gate names, commit hashes, captured- evidence file paths.
 - Cross-references to Issues.md / Fixed.md items by ATM-NNN (§11.4.54).
4. **HTML + PDF exports** travel with the markdown (identical mtimes within sync wrapper granularity).
5. **Auto-generation** — the canonical helper `scripts/testing/ generate_status_summary.sh <Status.md path>` produces both pages. The generator MUST be invoked from a sync wrapper (e.g. `scripts/testing/ sync_integration_status.sh` per §11.4.45) every time the source Status.md is updated.

Page-1 audience derivation. The generator maps domain → audience name via a project-controlled mapping table (e.g. `audio` → `Audio team` , `video` → `Video team` , `network` → `Network team` , `drm` → `DRM / streaming team`). Unmapped domains default to `Project team` . The mapping table lives in the project's generator config, NOT in the Constitution (project-specific data per §11.4.17).

Composition:

- §11.4.45 (Status.md per integration) — Status_Summary.md COMPLEMENTS, never replaces. Both files exist together.
- §11.4.12 + §11.4.53 (parity discipline) — Status_Summary follows the same regeneration / HTML+PDF / mtime-alignment pattern.
- §11.4.44 (revision header) — applies to Status_Summary.md.
- §11.4.23 (colorizer) — Status_Summary.html receives the same type/status visual cues if it embeds tracked-item references.
- §12.10 (CONTINUATION.md) — non-developer stakeholders read Status_Summary; engineers read Status.md + CONTINUATION.md + Issues.md.

Pre-build gates:

- CM-STATUS-SUMMARY-EXISTS-FOR-EVERY-STATUS — for every docs/**/Status.md, the sibling Status_Summary.md MUST exist (+ HTML + PDF) with $\text{mtime} \geq \text{Status.md mtime} - \text{tolerance}$.
- CM-STATUS-SUMMARY-TWO-AUDIENCE — every Status_Summary.md contains BOTH the Page 1 – For the <team> heading AND the Page 2 – For software engineers heading. Missing either FAILs.
- CM-STATUS-SUMMARY-REVISION-HEADER — §11.4.44 header present.
- CM-COVENANT-114-56-PROPAGATION — anchor literal across files.

Paired mutations (per §1.1):

- Delete a Status_Summary.md → CM-STATUS-SUMMARY-EXISTS-FOR-EVERY-STATUS FAILs.
- Remove Page 1 heading from a Status_Summary.md → CM-STATUS-SUMMARY-TWO-AUDIENCE FAILs.
- Strip §11.4.56 literal from constitution/CLAUDE.md → CM-COVENANT-114-56-PROPAGATION FAILs.

No escape hatch. No `--skip-status-summary`, `--engineer-only`, `--no-audience-split` flag. Plain-language stakeholder access is not negotiable — the project ships to humans, not just to engineers.

Classification: universal (per §11.4.17). Applies wherever §11.4.45 applies (every project that maintains domain-scoped Status.md docs).

2026 05 19

§ . . — README.md doc-link section +
revision metadata (User mandate,
- -)

Forensic anchor — verbatim user mandate (2026-05-19):

"Add a doc-link section to README.md — links to Issues + Issues_Summary + Fixed + Fixed_Summary + CONTINUATION + ALL Status docs + their exports. Each link shows revision + last-modified."

Why this anchor exists. Operators, AI agents, and external reviewers arriving at the project repo for the first time need a single discoverable entry point listing every canonical tracked- items + status document, along with each document's freshness (revision + last-modified). Without this surface, agents resort to `find docs -name '*.md'` which returns hundreds of files unranked. §11.4.57 makes the README.md the canonical index.

Operative rule. Every project's top-level README.md MUST contain a section titled Tracked-Items + Status Documents (or the project's localized equivalent — the section heading MUST contain the literal Tracked-Items for gate detection). The section MUST be a markdown table with columns: Document (human- readable name), Last modified (ISO 8601 UTC from the doc's §11.4.44 revision header), Revision (integer from same header), Markdown (relative link to .md), HTML (relative link to .html), PDF (relative link to .pdf).

The section MUST link to:

1. Issues.md , Issues_Summary.md (§11.4.12, §11.4.15, §11.4.16).
2. Fixed.md , Fixed_Summary.md (§11.4.19, §11.4.53).
3. CONTINUATION.md (§12.10).

4. **Every** `docs/**/*.Status.md` and its `Status_Summary.md` pair (§11.4.45, §11.4.56). Status docs are auto-discovered by the generator via `find docs -name 'Status.md' .`

Generator. `scripts/testing/update_readme_doc_links.sh` is the canonical helper. It:

1. Scans the canonical doc paths.
2. Extracts each doc's `Revision` + `Last modified` fields from its §11.4.44 header.
3. Renders the markdown table with current values.
4. Replaces the section between explicit markers
(`<!-- doc-link-section:begin -->` and `<!-- doc-link-section:end -->`) in `README.md`.

The wrapper **MUST** be invoked from `sync_issues_docs.sh` (so every `Issues.md` / `Fixed.md` update refreshes the `README` freshness data) AND from `sync_integration_status.sh` (so every `Status.md` update likewise refreshes `README`).

Composition:

- §11.4.12 + §11.4.19 + §11.4.53 (`Issues` / `Issues_Summary` / `Fixed` / `Fixed_Summary` parity) — `README` links to all four.
- §11.4.44 (revision header) — `README` pulls `Revision` + `Last modified` from each doc's header.
- §11.4.45 (`Status.md`) + §11.4.56 (`Status_Summary.md`) — `README` enumerates every Status pair.
- §12.10 (`CONTINUATION.md`) — explicitly linked from `README` so operators discover it on arrival.

Pre-build gates:

- `CM-README-DOC-LINK-SECTION-PRESENT` — `README.md` contains the literal `Tracked-Items` heading AND both section markers (`<!-- doc-link-section:begin -->` , `<!-- doc-link-section:end -->`).
- `CM-README-DOC-LINK-ROWS-COMPLETE` — every canonical doc (`Issues`, `Issues_Summary`, `Fixed`, `Fixed_Summary`, `CONTINUATION`, every `Status.md` + `Status_Summary.md` found under `docs/`) appears as a row inside the section.

- `CM-README-DOC-LINK-FRESHNESS` — the `Last modified` value in each row matches the corresponding doc's §11.4.44 header within sync-wrapper granularity (no row staler than the source doc by more than the sync wrapper's tolerance).
- `CM-COVENANT-114-57-PROPAGATION` — anchor literal across files.

Paired mutations (per §1.1):

- Strip the section markers from `README.md` → `CM-README-DOC-LINK-SECTION-PRESENT` FAILs.
- Remove the `CONTINUATION.md` row from the section → `CM-README-DOC-LINK-ROWS-COMPLETE` FAILs.
- Backdate a `Last modified` value in a row → `CM-README-DOC-LINK-FRESHNESS` FAILs.
- Strip §11.4.57 literal from `constitution/CLAUDE.md` → `CM-COVENANT-114-57-PROPAGATION` FAILs.

No escape hatch. No `--skip-readme-doc-links`, `--collapse-status-rows`, `--no-freshness-check` flag. The discipline exists because invisible docs are non-existent docs from a discoverability standpoint, and the §11.4 covenant specifically prohibits surfaces where claims (here: doc freshness) can drift unnoticed.

Classification: universal (per §11.4.17). Applies to every project that maintains an Issues / Fixed split or Status.md docs (i.e., to every project consuming this Constitution that has any of §11.4.15 / §11.4.45 / §11.4.53 in scope).

2026 05 19

§ . . — Parallel-development methodology (User mandate, — —)

Forensic anchor — verbatim user mandate (2026-05-19T~05:00Z MSK):

"We MUST DO one comprehensive research and planning in the background in parallel with current mainstream work: our current methodology of the development is very slow. It takes us days to fix a bunch of bugs, add some changes or features and all this to be tested, verified and shipped. From iteration to iteration we are slower and slower. We MUST CREATE adjusted improved version of working methodology where multiple workable items could be done in parallel, all come into the central (main) branch as they are done, then we at particular moment rebuild and reflash the System and in background full testing with validation and verification is done. Each parallel work on one workable (or more) item(s) must use parallel agents as much as possible! We MUST HAVE proper synchronization mechanism between parallel working routines so they do not cause conflicts, break features or create any other types of the problems! Document everything - the whole plan, whole methodology, working flow with in depth diagrams and graphs, all user guides and manuals and then add this all into our root (constitution Submodule) Constitution, CLAUDE.MD and AGENTS.MD! Make sure we start using this ASAP since we MUST increase efficiency and productivity a couple of times now! Writing in depth tests (all supported types of the tests) with Challenges and full HelixQA use is MANDATORY! Every test we execute besides executed with success MUST RESULT in proof that actual functionality being tested REALLY DOES WORK with NO BLUFF of any kind! Heavt enforcement of no-bluff / anti-bluff policy IS MANDATORY!"

Why this anchor exists. Sequential phase-chain working pattern (Phase 39.A → Phase 39.B → ... → Phase 39.DT , observed 2026-05-05 to 2026-05-19: 86 unique Phase 39.X sub-phases across 14 days = ~6/day sub-phase rotation, 224 commits / 14 days = ~16/day) serialises on five distinct bottlenecks: (B1)

`commit_all.sh flock at .git/.commit_all.lock` , (B2) single-thread sub-phase chaining via `docs/CONTINUATION.md` hand-offs, (B3) rebuild-blocking on every `REQUIRES_REBUILD` fix (§12.7 -j2 cap = 4–6 h rebuild + 30 min flash

- 60 min validate per cycle), (B4) single-thread `test_all_fixes.sh` sweep (~30–60 min × 8 phase-device combinations), (B5) subagent dispatch latency (one-at-a-time spawn + wait pattern). Net effect: single-item wall-clock 1–4 days where productive work is < 4 h. The User mandate requires the methodology to multiply throughput several-x.

Operative rule. Project work proceeds through the Parallel Work Unit (PWU) pipeline rather than sequential phase-chain. Each PWU is a self-contained workable item with mandatory components (ATM-NNN identifier per §11.4.54, Issues.md entry per §11.4.15+§11.4.16, file- scope manifest, §11.4.43 RED test, source patch, pre-build gate per §11.4.4(b) layer 1, post-flash test per §11.4.4(b) layer 3, paired §1.1 meta-test mutation, §11.4.4(b) layer 4 HelixQA Challenge bank entry, captured-evidence directory per §11.4.5+§11.4.52). PWUs execute through five pipeline stages: **Stage 1 DEVELOP** (parallel, N PWU agents in isolated worktrees), **Stage 2 MERGE** (serial, conductor processes one PWU at a time via `commit_all.sh flock + §11.4.41 4- step merge-first pipeline`), **Stage 3 REBUILD+FLASH** (parallel where possible: AOSP build inside §12.7 bounded scope, D3+D4 flash serialised on USB hardware), **Stage 4 VALIDATE** (parallel on D3+D4+meta-test+ coverage-audit), **Stage 5 SWEEP** (parallel: HelixQA Challenge bank, Issues → Fixed migration, README.md doc-link refresh per §11.4.57, HTML+PDF exports). Stage 1 of round N+1 starts WHILE Stages 4+5 of round N are still running — this is the primary throughput multiplier.

Synchronization mechanism. Four-layer lock hierarchy: **L1** parent serial flock `.git/.commit_all.lock` (held only during Stage 2, released between PWUs), **L2** per-submodule git operations (implicit via cascade), **L3** contention-path advisory locks at `qa-results/pwu-locks/<sha256(path)>.lock` for the 10 forbidden cross- PWU paths (CLAUDE.md/AGENTS.md/Constitution.md/ CONTINUATION.md/ Issues.md/Fixed.md/pre_build_verification.sh/ meta_test_false_positive_ proof.sh/build.sh/commit_all.sh+push_all.sh/ device.mk+BoardConfig.mk/ atmosphere-.sh/init.rc), **L4** per-PWU git worktree (zero cross- PWU file conflicts during Stage 1). Disjoint-scope PWUs (different test

files, different APKs, different drivers, different framework services) run fully parallel. Conflict detection at Stage 2 entry: `git fetch --all --prune --tags + git diff main...HEAD -- <scope>`

- cross-PWU scope overlap check. Overlap detected → PWU rejected to Stage 1 for rebase (per §11.4.41 step 2 integration-before-force-push).

Anti-bluff enforcement. Conductor REFUSES to merge a PWU lacking all four anti-bluff checks: **C1** §11.4.43 RED test captured-evidence (file exists + non-zero exit code in log + timestamp predates source patch — proves test catches regression), **C2** §1.1 paired meta-test mutation (applied + gate FAILs under mutation + restored + gate PASSes — proves gate is not a bluff gate), **C3** §11.4.50 deterministic- consistency (3 iterations for normal, 10 for cycle-validation, ALL identical exit codes AND identical evidence-hashes — eliminates §11.4.7 flake/transient/intermittent demotion path), **C4** §11.4.5 captured-evidence (audio: WAV with non-trivial RMS + fprobe channel count; video: screen recording with fprobe frame count > 0 + analyzer matched event; UI: uiautomator dump with under-test element state; or HelixQA `result.json` PASS verdict + positive-evidence chain). HelixQA coverage is MANDATORY for every user-visible PWU per User mandate — Challenge bank entry in `tools/helixqa/banks/atmosphere.yaml` referencing the PWU's ATM-NNN + dispatching to the on-device test + scoring PASS only on positive captured evidence. Metadata-only PASS / configuration- only PASS / absence-of-error PASS / grep-without-runtime-evidence PASS all REJECTED at the merge gate.

Agent orchestration. Four roles: **Conductor** (1, main thread) — issues ATM-NNN, dispatches PWU agents, runs Stage 2 merge serially, runs Stages 3-5, owns all contention-path edits, maintains `docs/CONTINUATION.md` §3 live snapshot. **PWU agent** (4–6 background) — one PWU each, Stage 1 work in isolated worktree, signals `qa-results/pwu-<atm>/READY_FOR_MERGE` on completion. **Validator** (4 background) — Stage 4 work in parallel on D3+D4+meta-test+coverage- audit. **Sweeper** (2 background) — Stage 5 work in parallel with next round's Stage 1. Max concurrent agent count ~12 at peak, bounded by §12.6 60% memory budget (~6 GB total ≤ 38 GB budget). Composition with §12 host-session safety: conductor refuses dispatch if `host_check_safety` fails; per §12.7 AOSP build still hard-capped at -j2 inside bounded scope; per §12.8 + §12.8b + §12.8c concurrency gates still block on stray gradle/kotlin daemons.

Composition map (mandatory):

- §11.4.4 — four-layer test coverage per PWU (pre-build + post-build + on-device + HelixQA bank).
- §11.4.5 — audio/video quality analysis comprehensiveness (PWU captured-evidence directory).
- §11.4.6 — no-guessing mandate (every PWU claim backed by captured forensic evidence or explicit `UNCONFIRMED:` tag).
- §11.4.7 — demotion-evidence rule (3/3 deterministic-consistency closes the demotion-bluff path).
- §11.4.9 — batch-source-fixes-before-rebuild (Stage 2 → Stage 3 batch threshold).
- §11.4.15 + §11.4.16 + §11.4.33 — PWU status + type vocabulary + type-aware closure terminal value.
- §11.4.19 — atomic Issues.md → Fixed.md migration on PWU closure.
- §11.4.41 — Pre-Force-Push Merge-First 4-step pipeline at Stage 2 entry.
- §11.4.42 — Iteration-discipline (PWU pipeline IS the iteration conductor; §11.4.58 binds steps 1-5 to PWU stages 1-5).
- §11.4.43 — TDD RED-first per-PWU enforced at merge time.
- §11.4.45 — Integration-status-doc maintenance (PWU's Status.md updated at every stage transition).
- §11.4.49 — Dual-approach testing (PWU MUST ship both UI-driven AND Intent/Broadcast-driven variants when applicable).
- §11.4.50 — Deterministic-consistency (3-iter normal / 10-iter cycle-validation merge-time enforcement).
- §11.4.52 — Autonomous-validation (PWU's on-device test MUST run without operator presence per Stage 4 parallel design).
- §11.4.54 — ATM-NNN ticket identifier (PWU identifier replaces Phase 39.X letters).
- §11.4.57 — README.md doc-link refresh at Stage 5.
- §12.6 — 60% memory-budget ceiling (caps concurrent agent count).
- §12.7 — AOSP `m -j` hard-cap at -j2 (Stage 3 build cap).
- §12.8 + §12.8b + §12.8c — concurrency hardening (Stage 3 preflight refuses on running gradle/kotlin/git-pack-objects/load-per-cpu spike).

- §12.10 — CONTINUATION.md maintenance (conductor owns the file by §3.4 forbidden cross-PWU rule).
- §9.2 — data safety (every history rewrite during Stage 2 merge-first step 2 backed by hardlinked `.git` backup).

Pre-build gates:

- `CM-PWU-LOCK-HIERARCHY` — verify lock hierarchy script exists at `scripts/testing/pwu_acquire_path_lock.sh` and `scripts/testing/pwu_release_path_lock.sh`, both executable, both reference the 10 forbidden cross-PWU paths from §3.4 of `docs/guides/PARALLEL_DEVELOPMENT_METHODOLOGY.md`.
- `CM-PWU-ANTI-BLUFF-COVERAGE` — for every PWU with status "Ready for merge" or beyond, assert all four anti-bluff checks (C1–C4) have captured-evidence files under `qa-results/pwu-<atm>/evidence/`. PWU lacking any of `red-test-output / mutation_patch / deterministic-consistency / captured-feature-proof` FAILs the gate.
- `CM-PWU-MERGE-QUEUE-DISCIPLINE` — scan `qa-results/pwu-queue/` for PWUs that bypassed the merge queue (commits on main without a corresponding `qa-results/pwu-queue/merged/<atm>` marker). Any bypass FAILs the gate.
- `CM-PWU-PARALLEL-AGENT-LIMIT` — assert no more than 6 concurrent background PWU agents (count `tmux/screen` sessions or `worktree` branches) — over-limit FAILs the gate (per §12.6 memory budget).
- `CM-COVENANT-114-58-PROPAGATION` — anchor literal §11.4.58 present in `constitution/CLAUDE.md` + `constitution/AGENTS.md` + parent `CLAUDE.md/AGENTS.md` + every owned-submodule `CLAUDE.md/AGENTS.md` (42 files total in the current consuming project's family).

Paired mutations (per §1.1):

- Move `scripts/testing/pwu_acquire_path_lock.sh` aside → `CM-PWU-LOCK-HIERARCHY` FAILs.
- Touch a PWU's `READY_FOR_MERGE` marker without `red-test-output.log` → `CM-PWU-ANTI-BLUFF-COVERAGE` FAILs.
- Land a commit on main without a corresponding `qa-results/pwu-queue/merged/<atm>` marker → `CM-PWU-MERGE-QUEUE-DISCIPLINE` FAILs.

- Spawn a 7th concurrent PWU agent → CM-PWU-PARALLEL-AGENT-LIMIT FAILs.
- Strip §11.4.58 literal from constitution/CLAUDE.md → CM-COVENANT-114-58-PROPAGATION FAILs.

No escape hatch. No `--skip-merge-queue` , `--allow-bypass` , `--no-anti-bluff-check` , `--unlimited-agents` , `--sequential-phase-chain-mode` flag exists in any pipeline helper. The discipline exists because (a) the User mandate explicitly requires the multiplier ASAP and (b) the §11.4 covenant specifically prohibits PASS-bluffs (which the four anti-bluff merge-time checks mechanically prevent). Operators who feel they need to bypass the pipeline for a specific high-urgency fix should split the work into a single trivial PWU and let the pipeline run normally — Stage 1 of a trivial PWU is < 30 min, Stage 2 is < 5 min, Stage 3 is the same rebuild cost they would pay anyway.

Classification: universal (per §11.4.17). Applies to every project consuming this Constitution that has multiple owned-submodules

- a sequential commit/push tool + a captured-evidence testing discipline. Project-specific implementations (agent dispatch helpers, worktree layout, batch thresholds) live in consumer-side `docs/guides/`
`PARALLEL_DEVELOPMENT_METHODOLOGY.md` (this Constitution defines the discipline; consumers implement the tooling).

Reference: parent `docs/guides/PARALLEL_DEVELOPMENT_METHODOLOGY.md` (full methodology with diagrams, templates, operator manual, and migration plan).

2026 05 19

§ . . — README always-sync mandate
(User mandate, — —)

Forensic anchor — direct user mandate (verbatim, 2026-05-19):

"fully review and update our main README document. Some points are not valid anymore, some are missing. Make sure main README is among documents we MUST ALWAYS keep updated and in Sync with the projects and other documentation! Make sure we always export it (on every update) into PDF and HTML. This mandatory rules/constraints MUST BE all added into the root (constitution Submodule) Constitution, AGENT.MD and CLAUDE.MD!"

`README.md` at the project root is a §11.4.12-class always-sync document. It MUST be:

1. **Reviewed and updated whenever** new docs / integrations / Status.md entries appear, new submodules land, applied-fixes count changes, or canonical paths shift.
2. **Kept in lockstep** with `docs/CONTINUATION.md` (§12.10) and the Issues / Issues_Summary / Fixed / Fixed_Summary doc set (§11.4.12, §11.4.53) — same always-sync discipline.
3. **Exported to .html and .pdf on every update.** New helper `scripts/testing/sync_readme_export.sh` performs the pandoc + weasyprint export; it is auto-invoked by `scripts/testing/sync_issues_docs.sh` so a single doc-sync run refreshes the entire doc surface (Issues / Issues_Summary / Fixed / Fixed_Summary / CONTINUATION / README — md + html + pdf).
4. **Carry a §11.4.44 revision header** (Revision N + Last modified ISO timestamp) at the top of the file.
5. **Contain a Documentation Map section** linking to every Status.md
 - Status_Summary.md + spec + plan + guide + script-companion doc + changelog + the constitution submodule, plus a per-audience navigation (end user / developer / QA / agent).
6. **Self-contained** — no hyperlinks to ephemeral external systems as the only source of truth.

Stale exports are §11.4.59 violations regardless of whether the underlying `README.md` is correct — an operator (or future agent) reading the HTML or PDF gets a divergent view, and the §12.10 CONTINUATION resumption guarantee silently breaks. Same discipline as §11.4.12 Issues_Summary and §11.4.53 Fixed_Summary applied to README.md.

Captured-evidence enforcement. Pre-build gate `CM-README-EXPORT-SYNC` locks four invariants: (a) `README.md` exists, (b) `README.html` exists, (c) `README.html` mtime $\geq$ `README.md` mtime, (d) `README.pdf` mtime $\geq$ `README.md` mtime (skipped gracefully if weasyprint is unavailable on the host but enforced when the PDF is present). Paired meta-test mutation backdates `README.html` + `README.pdf` so the source is newer $\rightarrow$ gate FAILs.

Composition. Composes with §11.4.12 (Issues_Summary parity), §11.4.18 (script-companion docs), §11.4.23 (visual-cue HTML colorizer not applicable to README's narrative format), §11.4.44 (revision header on every Status.md and now README.md), §11.4.45 (Status.md integration), §11.4.53 (Fixed_Summary parity — README is the canonical sibling at the project-overview layer), §11.4.56 (Status_Summary parity), §11.4.57 (README.md is the canonical index per Phase 39.EW+1), §12.10 (CONTINUATION.md resumption guarantee depends on README's Documentation Map being current).

No escape hatch. §11.4.59 has NO operator-facing override flag. No `--skip-readme-sync`, `--no-readme-export`, `--readme-stale-OK` flag exists. The discipline exists because README.md is the first file any new agent / operator / contributor reads — letting it diverge from reality is the exact §11.4 PASS-bluff pattern but at the project-overview layer.

Classification: universal (per §11.4.17). Applies to every project consuming this Constitution. The doc-sync helper, pandoc + weasyprint dependencies, and gate name are universal; the specific Documentation Map contents are consumer-side per-project.

Canonical authority: constitution submodule `Constitution.md` §11.4.59.

Non-compliance is a release blocker regardless of context.

§ . . — Documentation always-sync composite covenant (User mandate, — —)

Forensic anchor — verbatim user mandate (2026-05-19 ~09:00Z):

"Double check if all documents are properly tied with our root Constitution, CLAUDE.MD and AGENTS.MD so they are always up to date, always in sync and exported into PDF and HTML! ... Issues, Issues_Summary, Fixed, Fixed_Summary, Continuation, Status and Status_Summary for all contexts (areas) — THEY ALL MUST BE REGULARLY UPDATED, IN SYNC AND CONSISTENT without giving at any moment false picture about the state of the project or particular area(s) of it!"

The covenant. Eight documentation classes constitute the project's living state surface. They MUST be in sync at all times across markdown + HTML + PDF artefacts. Per-class anchors §11.4.12 / §11.4.44 / §11.4.45 / §11.4.53 / §11.4.56 / §11.4.57 / §11.4.59 / §12.10 each govern one class; §11.4.60 binds them via a single mechanically-enforced composite invariant so the failure mode "one per-class gate silently disabled while operator reads a divergent HTML/PDF" is structurally impossible.

Eight bound doc classes (artefact triple — .md + .html + .pdf):

1. docs/Issues.md — open-item tracker (§11.4.12 + §11.4.15 + §11.4.16)
2. docs/Issues_Summary.md — auto-generated short form (§11.4.12)
3. docs/Fixed.md — closed-item archive (§11.4.53)
4. docs/Fixed_Summary.md — auto-generated short form (§11.4.53)
5. docs/CONTINUATION.md — resumption handoff (§12.10)
6. README.md — project overview + doc-link index (§11.4.57 + §11.4.59)
7. docs/**/Status.md (every domain-scoped instance) — per-integration status (§11.4.45)
8. docs/**/Status_Summary.md (every domain-scoped instance) — auto-generated short form (§11.4.56)

Mandatory composite gate. `CM-DOCS-COMPOSITE-SYNC` (`pre-build, device/rockchip/rk3588/tests/pre_build_verification.sh`):

- For every doc-class instance with `.md` present, verify `.html` `mtime` $\geq$ `.md` `mtime` AND `.pdf` `mtime` $\geq$ `.md` `mtime`.
- For every `Status.md` instance, also verify the sibling `Status_Summary.md` `mtime` $\geq$ `Status.md` `mtime` (if the summary exists).
- Walks `docs/` recursively for `Status.md` — the Status fleet is not enumerated statically because new integrations land per phase.

Composite gate FAILs the build if ANY single instance fails. The per-class gates (`CM-ISSUES-SUMMARY-SYNC` , `CM-DOCS-EXPORT-SYNC` , `CM-FIXED-SUMMARY-SYNC` , `CM-CONTINUATION-DOC-INSYNC` , `CM-README-EXPORT-SYNC`) remain in place — §11.4.60 is the belt-and-suspenders layer that catches what they miss.

Paired mutation (per §1.1). `meta_test_false_positive_proof.sh` backdates `docs/Issues.html` to year 2000 (`touch -t 200001010000.00`) and asserts `CM-DOCS-COMPOSITE-SYNC` FAILs.

Revision-header coupling. Per §11.4.44 every one of the 8 doc classes carries the `**Revision:** N + **Last modified:** ISO 8601 UTC` header directly under the H1. `CM-DOCS-COMPOSITE-SYNC` does NOT re-check the revision header (already covered by `CM-DOC-REVISION-HEADER-PRESENT`) — composite gate is artefact-mtime only, by design narrow.

Auto-sync wrapper coupling. Per §11.4.12 every edit to an Issues/Fixed-class doc flows through `sync_issues_docs.sh` . Per §11.4.45 every edit to a `Status.md` flows through `sync_integration_status.sh <Status.md>` . Per §11.4.59 every edit to `README.md` flows through `sync_readme_export.sh` . §11.4.60 does NOT change the wrapper contract — it only catches the failure mode where an operator (or AI agent) bypassed the wrapper and committed markdown without regenerating exports.

No escape hatch. No `--skip-composite-doc-sync` , `--allow-stale-html` , `--summary-not-applicable` flag exists. The covenant exists for the operator's own protection — the moment a single per-class gate is bypassed, divergence accumulates silently until release-time forensics finds it (the exact §11.4 PASS-bluff pattern the canonical covenant was authored to eliminate).

Composes with §11.4.12 (Issues_Summary), §11.4.15 (Status field), §11.4.16 (Type field), §11.4.19 (atomic Issues → Fixed migration), §11.4.23 (colorizer), §11.4.33 (type-aware closure vocabulary), §11.4.44 (revision header), §11.4.45 (Status.md), §11.4.53 (Fixed_Summary), §11.4.56 (Status_Summary), §11.4.57 (README doc-link), §11.4.59 (README always-sync), §12.10 (CONTINUATION maintenance).

Propagation. Pre-build gate `CM-COVENANT-114-60-PROPAGATION` enforces the §11.4.60 anchor literal in every CLAUDE.md / AGENTS.md across parent + 10 owned submodules + HelixQA dependencies (same shape as existing §11.4.5X propagation gates). The propagation gate is OPTIONAL for Phase 39.FC's initial landing — only the composite gate + its paired mutation are mandatory; propagation gates trail per established phase pattern.

Classification: universal (per §11.4.17). Applies to every project consuming this Constitution that maintains any of the 8 doc classes. The composite gate logic is universal; the specific doc-class instance paths are consumer-side per-project.

Canonical authority: constitution submodule `Constitution.md` §11.4.60.

Non-compliance is a release blocker regardless of context.

2026 05 19

§ . . — Mandatory Markdown metadata
table + structured-doc ToC (User mandate,
— —)

Forensic anchor — direct user mandate (verbatim, 2026-05-19):

"For every Markdown document which contains structured content (with headings / sections and sub-sections) make sure that every time we apply change to the structure, table of contents on the top of the document is created or updated! This is MANDATORY for every structured MARKDOWN document. Automatically its PDF and HTML versions MUST BE (re)generated!"

"Introduce ... revision number, date and time of creation, date and time of last modification, other useful information we have in documents such as Issues, Issues_Summary, Fixed, Fixed_Summary, Status, Status_Summary, Continuation and similar. ... make this mandatory for EVERY Markdown document from now on, update root constitution Submodule with these changes and commit and push it all to all upstreams."

Scope of §11.4.61. This clause adds two universal disciplines that sit alongside (NOT replacing) the §11.4.65 universal export mandate. §11.4.65 governs the existence and freshness of `.html` + `.pdf` siblings; §11.4.61 governs the *content* of the Markdown sources: how revision/status metadata is displayed at the top, and how the document's navigation map (Table of contents) is kept in lockstep with the heading structure. The two clauses compose — §11.4.65 guarantees readers see *something current*; §11.4.61 guarantees what they see communicates its own provenance and is internally navigable.

A. Canonical metadata table (supersedes §11.4.44 format). The §11.4.44 bold-line revision header is superseded by a **Markdown table** placed immediately after the document's H1 title (and any blank line that follows). Pre-existing bold-line headers remain historically valid but MUST migrate to the canonical table at the next substantive edit. The canonical table has these MANDATORY rows:

Field	Value
Revision	positive integer; starts at 1, monotonic, never reset
Created	ISO 8601 date — first commit touching the file
Last modified	ISO 8601 date — most recent commit touching the file
Status	one of <code>active</code> / <code>draft</code> / <code>deprecated</code> / <code>superseded</code>

ENCOURAGED rows (REQUIRED when the document tracks workable items or has a known continuation; render with `–` / `none` when N/A):

Field	Value
Status summary	one-line current-state hook
Issues	comma-separated workable-item IDs
Issues summary	one-line summary of open issues
Fixed	comma-separated workable-item IDs of resolved items
Fixed summary	one-line summary of fixes
Continuation	next action or linked CONTINUATION.md anchor

Additional rows MAY be added for project-specific metadata as long as the MANDATORY rows are present.

B. Structured-document Table of contents. Any tracked `*.md` with **two or more H2 sections** ("structured content") MUST include a `## Table of contents` section immediately after the metadata table. The ToC MUST list every H2 and H3 in document order with anchor links and MUST be regenerated on every structural change (heading added / removed / renamed / reordered). A stale ToC is a §11.4 PASS-bluff: operators reading the rendered Markdown see a navigation map that no longer matches the body.

Scope (IN-scope under §11.4.61). Every tracked `*.md` in the §11.4.65 INCLUDED set (project-root `*.md` , `docs/**/*.*md` , `scripts/**/*.*md` companion docs, owned-submodule trees, `constitution/**/*.*md` , owned

HelixQA dependencies). The §11.4.44 exclusions for `CLAUDE.md` / `AGENTS.md` / `README.md` are SUPERSEDED here too — explicit metadata in the file itself is preferred over indirect tracking via a separate `VERSION` file because operators reading the rendered document see the freshness directly.

Narrow exemptions (NOT IN-scope). Same EXCLUDED set as §11.4.65 (`external/**` , `prebuilts/**` , `out/**` , `build/**` , application- internal `*.md` shipped with source code, non-owned third-party submodule trees). Additional exemptions for the structural rules: `LICENSE` , `LICENSE.md` , `NOTICE` , `VERSION` , `OWNERS` , machine- generated `CHANGELOG.md` .

PDF + HTML siblings. Deferred to §11.4.65. The metadata table and ToC are *content* requirements; their export-freshness guarantee lives in §11.4.65's `CM-UNIVERSAL-MARKDOWN-EXPORT-SYNC` gate. A §11.4.61 metadata or ToC change is a `.md` modification, which transitively triggers §11.4.65 regeneration — the two clauses compose without duplicate enforcement.

Anti-bluff captured-evidence gates (planned):

- `CM-MD-METADATA-PRESENT` — walks every tracked `*.md` (minus exemptions) and asserts (a) H1 present, (b) within 25 lines below H1 a Markdown table contains all four MANDATORY rows (`| Revision |` , `| Created |` , `| Last modified |` , `| Status |`).
- `CM-MD-TOC-PARITY` — for every `*.md` with ≥ 2 H2 sections, asserts a `## Table of contents` is present and its entries' anchor slugs match the document's live H2/H3 set in order.

Paired §1.1 mutation tests (planned):

- Strip the `| Revision |` row from one `*.md` → `CM-MD-METADATA-PRESENT` FAILs.
- Rename one H2 without updating ToC → `CM-MD-TOC-PARITY` FAILs.

Composition with §11.4.44. §11.4.44 remains the authoritative clause for the *fact* that documents must carry revision metadata; §11.4.61 changes the *format* (bold-lines → table) and *scope* (`docs/**/*.*md` → every tracked `*.md` per §11.4.65 INCLUDED set). Consumer projects with existing gates that grep for `**Revision:**` MUST update those gates to also accept the canonical table form. Migration period: 30 days from §11.4.61 landing.

Composition with §11.4.45 / §11.4.59 / §11.4.60 / §11.4.65. The per-class always-sync clauses (Status.md §11.4.45, README §11.4.59, composite §11.4.60) and the universal export covenant §11.4.65 govern *artefact existence and freshness*; §11.4.61 governs *content discipline inside the source .md* . The clauses are orthogonal layers of the same anti-bluff posture.

No escape hatch. No `--skip-md-metadata` , `--no-metadata-table` , `--toc-stale-OK` , `--allow-missing-revision` flag exists. Divergent metadata and stale ToCs are §11.4 PASS-bluff patterns at the doc-surface layer.

Classification: universal (per §11.4.17). Applies to every project consuming this Constitution. The canonical table format and ToC mandate are universal; the per-project tracker IDs (e.g. `HRD-`) and specific Issues/Fixed-class doc paths are consumer-side per-project.

Canonical authority: constitution submodule `Constitution.md` §11.4.61.

Non-compliance is a release blocker regardless of context.

2026 05 19

§ . . – Workable-items procedure docs
as single source of truth (User mandate,
– –)

Forensic anchor — verbatim user mandate (2026-05-19 ~10:30Z):

"To make workable items tracking and creation through all of the mentioned mandatory documents we MUST create proper procedure document which will be named in Constitution, AGENTS.md and CLAUDE.md as single source of truth for this procedure! Make sure it is created under: docs/procedures/issues/Creation.md which will be always exported to PDF and HTML whenever it is updated. We MUST create procedure documents for Updating.md, Resolution.md, Reopening.md and other workable items actions that we have. ... Everything done on project - new features, changes, tasks, bug fixes, investigation, ordinary tasks (writing documentation for example) MUST ALL go through the workable items flow and proper procedures!"

Every workable-item action — opening, updating, closing, reopening, migrating across Issues.md ↔ Fixed.md — MUST follow the canonical procedure document at docs/procedures/issues/<Action>.md . The five procedure docs are MANDATORY references; no agent or operator may invent ad-hoc procedures. The closed-set:

Procedure doc	Covers
Creation.md	Opening a new workable item (Bug / Feature / Task)
Updating.md	Non-closure, non-reopen edits — status transitions, evidence append, type re-classification
Resolution.md	Atomic Issues.md → Fixed.md close (Bug → Fixed, Feature → Implemented, Task → Completed per §11.4.33)
Reopening.md	Atomic Fixed.md → Issues.md reopen with §11.4.34 source attribution + §11.4.55 reopens-history
Migration.md	Bidirectional atomic-move mechanics + sync_issues_docs.sh internals (consumed by Resolution + Reopening)

Each procedure doc carries the §11.4.44 revision header + HTML + PDF exports synchronised to its .md source. New helper scripts/testing/sync_procedure_docs_export.sh walks the procedure- docs directory and emits matching .html + .pdf for each .md ; invoked as the final stage of scripts/testing/sync_issues_docs.sh so a single operator invocation refreshes every doc-side artifact.

Universality of scope. All work — new features, behavioural changes, tasks (refactor / doc / infra / gate / audit / cleanup), bug fixes, investigation, documentation work — MUST flow through these procedures. The "I'll just tweak this one file" escape hatch is forbidden because it produces silent doc drift and reopens the §11.4 covenant's PASS-bluff failure mode at the process-tracking layer.

Composes with §11.4.6 (no-guessing — every procedure step is mechanical, no likely / seems), §11.4.7 (demotion-evidence — Reopening.md is the canonical implementation), §11.4.11 (file-layout — procedure docs live in docs/procedures/issues/ , qa-results forensics live in qa-results/), §11.4.12 (Issues_Summary

sync — all five procedures invoke `sync_issues_docs.sh` whose pipeline is documented in `Migration.md`), §11.4.15 (Status closed-set — all five procedures cite the same closed-set), §11.4.16 (Type closed-set), §11.4.19 (atomic Issues ↔ Fixed move — `Migration.md` is the implementation), §11.4.23 (visual cue + grouping colorization — invoked by sync wrapper, no procedure-side bypass), §11.4.33 (type- aware closure vocabulary — `Resolution.md` enforces), §11.4.34 (Reopened-source attribution — `Reopening.md` enforces), §11.4.44 (revision header on every procedure doc), §11.4.53 (Fixed_Summary parity — `Migration.md` documents the pipeline), §11.4.54 (ATM-NNN identifier — `Creation.md` allocates), §11.4.55 (Reopens-history per- item doc — `Reopening.md` authors), §11.4.60 (documentation always- sync composite — procedure docs are an additional doc class binding to the composite covenant).

Pre-build gate `CM-PROCEDURES-DOCS-PRESENT` checks (1) all 5 procedure docs exist at `docs/procedures/issues/{Creation,Updating,Resolution,Reopening,Migration}.md`, (2) each carries the §11.4.44 revision header, (3) each has matching `.html` + `.pdf` exports with `mtime` ≥ the `.md` `mtime`, (4) `scripts/testing/sync_procedure_docs_export.sh` exists + is executable, (5) `scripts/testing/sync_issues_docs.sh` invokes `sync_procedure_docs_export.sh`. Paired mutation: rename one procedure doc → gate FAILs.

Propagation gate `CM-COVENANT-114-63-PROPAGATION` enforces this anchor literal in every `CLAUDE.md` / `AGENTS.md` across parent + 10 owned submodules + nested submodules + HelixQA dependencies. Paired mutation strips the anchor literal → gate FAILs.

No escape hatch — no `--ad-hoc-procedure`, `--skip-procedure-doc`, `--procedure-not-applicable` flag exists. Inventing an alternate procedure for a "small" change is the exact failure mode this anchor closes.

Classification: universal (per §11.4.17). The closed-set of five procedure docs and the sync-wrapper composition are universal; the specific test-name / fix-name examples within each procedure doc may be consumer-side per-project.

Canonical authority: constitution submodule `Constitution.md` §11.4.63.

Non-compliance is a release blocker regardless of context.

§ . . — Universal Markdown export mandate (User mandate, — —)

Forensic anchor — direct user mandate (verbatim, 2026-05-19):

"Any markdown document inside the project and which is not part of the applications or services source code **MUST BE** exported (be available) in PDF and HTML! Any already existing Markdown document that fulfills this condition and which does not have HTML or PDF at all or it is not in sync with it **MUST HAVE** (re)generated PDF and HTML version! Every time when Markdown document (file) is modified, its proper HTML and PDF versions **MUST BE** regenerated. Markdown documents **MUST BE** at all times in sync with PDF and HTML versions!"

The covenant generalizes the per-class always-sync discipline of §11.4.12 / §11.4.18 / §11.4.44 / §11.4.45 / §11.4.53 / §11.4.56 / §11.4.57 / §11.4.59 / §11.4.60 / §11.4.63 / §11.4.64 to **every Markdown document in the project that is not part of an application or service's source-code tree**. Per-class anchors govern specific files (Issues, Fixed, Status, README, etc.); §11.4.65 is the catch- all: any *other* `.md` file that documents the project — guides, research notes, plans, hardware-ID notes, script-companion docs, changelogs, procedure docs, README files inside owned submodules — **MUST** also have `.html` + `.pdf` siblings, all three artefacts in sync at all times. Stale exports of *any* documentation surface are the exact "operator reads divergent HTML/PDF while `.md` is correct" PASS-bluff §11.4 forbids — only the per-class enforcement so far prevented it for specifically-named docs; §11.4.65 closes the long tail.

Scope (closed-set):

INCLUDED:

- Project root `*.md` (README.md, CLAUDE.md, AGENTS.md, CONTRIBUTING.md, etc.)
- `docs/**/*.*md` (guides, research, plans, changelogs, procedures, hardware notes)
- `scripts/**/*.*md` (script-companion docs in documentation format, NOT shebang scripts)

- Owned-submodule trees at `device/rockchip/atmosphere/<submodule>/` : top-level README.md / CLAUDE.md / AGENTS.md / CHANGELOG.md and any `docs/**/* .md`
- `constitution/**/* .md` (the canonical-root submodule)
- `tools/helixqa/<helixqa-submodule>/` top-level README.md / CLAUDE.md / AGENTS.md and any `docs/**/* .md` (owned HelixQA dependencies)

EXCLUDED (NOT subject to §11.4.65 — these are application/service source code or third-party trees we do not own):

- `external/**` (AOSP-mainline / upstream-mirror source trees)
- `prebuilts/**` (binary prebuilts + their bundled docs)
- `packages/modules/**` (AOSP mainline modules)
- `kernel-5.10/**` (kernel source tree — has its own upstream docs)
- `out/**` (build output)
- `build/**` (AOSP build system)
- Application / service source-code trees (e.g. Java/Kotlin module-internal `*.md` shipped with library code, README files inside source-only third-party gradle module directories)
- Any third-party submodule NOT in the owned-submodule set (presenter, vlc-player, nova-player, mpv-player, gramophone-player, rhythm-player, strep-player, smarttube-player, torrservice, lampa) or HelixQA-owned set (Challenges, Containers, DocProcessor, LLMOrchestrator, LLMPProvider, VisionEngine)

Mandatory protections (ALL must hold):

1. **Every INCLUDED .md file has .html and .pdf siblings.** A missing export is a §11.4.65 violation regardless of when the markdown was last touched.
2. **.html and .pdf mtime ≥ .md mtime** (within the same sync-wrapper invocation granularity). Stale exports are violations even if the .md itself is correct.
3. **Every modification triggers regeneration.** Whether via direct sync-wrapper invocation (`sync_all_markdown_exports.sh`), a git pre-commit hook auto-regenerating on staged `.md` changes, or the existing per-class wrappers

(§11.4.12 `sync_issues_docs.sh`, §11.4.18 `script-docs sync`, §11.4.59 `sync_readme_export.sh`, etc.) all of which delegate the universal export path to the same canonical helper.

4. **Pre-build gate enforces parity.** `CM-UNIVERSAL-MARKDOWN-EXPORT-SYNC`

walks the INCLUDED scope, verifies every `.md` has `.html`

- `.pdf` siblings with `mtime ≥ .md mtime`, and FAILs the build if any are missing or stale.

5. **No escape hatch.** No `--skip-md-exports`, `--no-pdf-only`,

`--md-export-not-applicable`, `--application-internal-doc` flag exists for files inside the INCLUDED scope. The EXCLUDED scope is the only legitimate path to opt out, and it is closed-set.

Canonical helper. `scripts/testing/sync_all_markdown_exports.sh` walks the INCLUDED scope, invokes pandoc (HTML) + weasyprint (PDF) with `timeout 60` each (graceful per-file degradation), uses `docs/_progress-style.css` for visual consistency, and supports `--check-only` (exit nonzero if any out-of-sync, print list) and `--regenerate-all` (force) modes. Caps at 500 candidates with explicit abort+list if the scope is unexpectedly large (e.g. after a new submodule lands without scope-update review). Idempotent.

Composition. Composes with §11.4.12 (Issues_Summary sync — universal helper is the back-end), §11.4.18 (script-companion docs), §11.4.23 (HTML colorizer continues to post-process the always-sync docs for visual cues), §11.4.44 (revision header on every `.md` the universal helper exports), §11.4.45 (Status.md auto-sync), §11.4.53 (Fixed_Summary), §11.4.59 (README always-sync — the universal helper covers the long tail of every *other* root-level `.md` too), §11.4.60 (composite always-sync covenant — §11.4.65 is the structural generalization), §11.4.63 (procedure docs — already always-sync; §11.4.65 adds catch-all for every doc class not yet anchored), §11.4.64 (topic discoverability — summary docs the universal helper exports continue to be post-processed by `inject_topic_index.py`).

Pre-build gates:

- `CM-UNIVERSAL-MARKDOWN-EXPORT-SYNC` — invariants: (a) `scripts/testing/sync_all_markdown_exports.sh` exists + executable, (b) running it in `--check-only` mode returns 0 (no out-of-sync file in the INCLUDED scope).

- `CM-COVENANT-114-65-PROPAGATION` — anchor literal `11.4.65` in every `CLAUDE.md` / `AGENTS.md` across canonical-root + parent + 10 owned submodules + nested submodules + HelixQA dependencies (42-file fleet — same propagation set as §11.4.58).

Paired meta-test mutations: one strips the `CM-UNIVERSAL-MARKDOWN-EXPORT-SYNC` gate's enforcement literal; one strips the §11.4.65 anchor literal from a sentinel propagation file. Both mutations assert the corresponding gate FAILs when applied.

Classification: universal (per §11.4.17). The mandate to export every project-doc Markdown to HTML+PDF is universal across every project that consumes this Constitution; the specific INCLUDED / EXCLUDED scope path lists are consumer-side per-project (since each consumer has different submodule topology, build trees, and source- code roots).

Canonical authority: constitution submodule `Constitution.md` §11.4.65.

Non-compliance is a release blocker regardless of context.

2026 05 19

§ . . — Blocker-resolution interactive-clarification mandate (User mandate, - -)

Forensic anchor — direct user mandate (verbatim, 2026-05-19):

"If any blockers which can be resolved with interactive response ever happen again, perform in depth research on options doable by your side and how much inputs from us you really need, then create options and present them to us. After we answer, preferrably you will be unblocked and be able to continue work on blocked items. Let us make this main approach when such situations (blockers) do happen!"

When any task is blocked (waiting on operator decision, hardware access, external-service authorization, ambiguous scope, or any input the agent cannot mechanically derive), the agent **MUST** follow this five-step discipline before idling or asking free-form questions:

1. **In-depth research on agent-side options.** Identify the maximum scope that can land unilaterally without operator input — research the codebase, upstream documentation, existing tooling, the current state, what files/components are involved, what tests/gates already exist. Surface every agent-actionable path even if it requires a small operator confirmation.
2. **Calculate minimum-viable operator input.** Reduce the question set to the smallest closed set of operator-only decisions (preference, authorization, scope-bounding, hardware availability, schedule). Anything that can be researched **MUST NOT** be asked.
3. **Construct 2–4 mutually-exclusive options.** Each option carries (a) a short label, (b) a one-line trade-off description, (c) an explicit statement of what the agent will do *after that answer*. One option marked "Recommended" with rationale. Options **MUST** genuinely differ — restating one option as three close variants is itself a §11.4.66 violation.
4. **Present via the platform's interactive question mechanism.** On Claude Code that is `AskUserQuestion` (max 4 questions, each with 2–4 options; supports multi-select for non-mutually-exclusive sets). Other platforms (Copilot CLI, Codex, Gemini CLI) use their equivalents. **NEVER** inline free-text "what would you like?" when interactive options would do — that wastes operator attention.
5. **After the operator answers, resume work without additional round-trips.** Every option's promised action **MUST** be sufficient to unblock — if a follow-up clarifying question is needed, the prior options were insufficiently researched and that's itself a §11.4.66 violation. The contract is: ask once, unblock, continue.

Composes with:

- §11.4.6 (no-guessing — interactive options replace agent guessing about operator preference)

- §11.4.7 (demotion-evidence — operator preference is captured- evidence for direction changes)
- §11.4.40 (full-suite retest — uninterrupted work post-answer means the test cycle resumes cleanly without context drift)
- §11.4.41 (merge-first — when operators make conflicting choices across sessions, the merge-first discipline preserves both)
- §11.4.42 (iteration-discipline — interactive-question loop is built into the per-cycle methodology)
- §11.4.52 (autonomous-validation — most blockers can be researched to autonomous-path level, leaving only true operator preference to interactive-ask)

Mandatory protections:

- **No silent waiting.** If blocked, ask. If genuinely no operator input is needed, do not ask — proceed.
- **No bulk-text questions when interactive options would do.** Free-text prompts are reserved for truly open-ended ideation, never for closed-set decisions.
- **Each interactive question minimizes operator cost** — the operator should be able to answer all questions in under 30 seconds total.
- **Recommended option is highlighted explicitly** — the operator should know what the agent would default to if no answer comes.
- **Out-of-band channels (slack, voice, email) NEVER substitute** — the interactive mechanism creates an audit trail that free-text channels do not.

Pre-build gate `CM-COVENANT-114-66-PROPAGATION` enforces the anchor literal is present across the 42-file consumer fleet (parent CLAUDE.md / AGENTS.md + 10 owned-submodule pairs + HelixQA-owned pairs). Paired meta-test mutation strips the literal → gate FAILs. No escape hatch — no `--skip-ask` , `--silent-wait` , `--free-form-only` flag is permitted.

Canonical authority: constitution submodule `Constitution.md` §11.4.66.

Non-compliance is a release blocker regardless of context.

\$. . – Shell-script target-shell-parseability mandate (User mandate, – –)

Forensic anchor — direct user mandate (verbatim, 2026-05-19):

"any issue we spot must be fixed, bash scripts as well if they are broken!"

"Make sure that this is mandatory rule!"

Forensic incident — Phase 39.FJ.67 D3 post-flash cycle block: `device/rockchip/rk3588/tests/test_all_fixes.sh:114` invoked on Android via `sh script.sh` used bash-only process substitution `exec > >(tee -a "$file") 2>&1`. Android's `/system/bin/sh` is mksh, which parses the *entire* script BEFORE executing it — so the runtime guard `if [ -n "${BASH_VERSION:-}" ]` could not save the script: mksh rejected `>( ... )` at parse time and the test cycle failed to launch on D3. The runtime guard correctly tried to gate the unsafe path, but mksh's compile-time parser made the guard useless. Fixed by wrapping the offending `exec` in `eval` so the parser sees only a string and the parse step defers to runtime where the guard can take effect.

This is the canonical class of script-bug FAIL-bluff: a test cycle PASSES on the developer host (bash) and FAILS at parse time on the target (mksh), producing a FAIL-bluff that masks every product defect the cycle would have caught. §11.4.1 already forbids script-bug FAILs. §11.4.67 mechanically prevents this specific subclass.

The mandate. Every shell script that may be invoked under a target shell other than the one in its shebang MUST parse cleanly under that target shell.

Specifically:

1. **Closed-set scope.** Every tracked `.sh` file under `device/rockchip/rk3588/tests/`, `scripts/`, and `scripts/testing/` (and equivalent paths in owned submodules) is in scope. Explicitly OUT of scope: `external/`, `prebuilts/`, `packages/modules/`, `kernel-5.10/`, `out/`, `build/`, `scripts/legacy/` (legacy quarantine), and any third-party vendored shell directory the project does not maintain.

2. **Mandatory parseability invariant.** Every in-scope script MUST parse under POSIX `sh -n` on the developer host. POSIX `sh` is the closest portable parser to Android `mksh` — passing `sh -n` on the host gives high confidence the script parses on `mksh`.
3. **Bash-only syntax requires deferral.** Process substitution `>(…) / <(…)` , `[[ ]]` , `<<<` here-strings, indexed arrays `arr=()` , associative arrays `declare -A` , `${var^^}` / `${var,,}` case-fold expansions, `${var/pattern/replacement}` with bash-only operators, `coproc` , `function name { … }` (vs POSIX `name() { … }`), and any other bash-only construct MUST EITHER be wrapped in `eval '…bash-only string…'` so the parser sees only a string OR be guarded by sourcing the file only on bash- detected hosts (`[ -n "${BASH_VERSION:-}" ]` BEFORE the script loads — runtime guards inside a single `mksh`-parsed script do not work).
4. **Honest shebangs.** A script that contains bash-only constructs without `eval -deferral` MUST declare `#!/bin/bash` (or equivalent) AND its callers MUST invoke it as `bash script.sh` rather than `sh script.sh` . A script with `#!/system/bin/sh` or `#!/bin/sh` shebang MUST be POSIX-clean — no bash-only syntax, no `local` keyword without an `mksh` fallback, no `read -p` , no `echo -e` without portable equivalent.
5. **eval -deferral is the safe pattern.** When a bash-only construct is genuinely needed inside a script that may be parsed by `mksh` (e.g. `sh script.sh` from a callsite outside our control), wrap the construct in single-quoted `eval : eval 'exec > >(tee -a "$f") 2>&1'` . The parser sees a string; the runtime guard around the `eval` decides whether to execute.
6. **Sourced-function exec redirection-scoping (extension, research-derived, 2026-07-23).** In bash-family shells, `exec` WITHOUT a command applies EVERY redirection on the line to the CURRENT SHELL, PERMANENTLY. Inside a standalone script that is a self-contained choice; inside a SOURCED / library function — code that runs in the operator's interactive shell or a caller's long-lived shell — it is a host-state mutation. Forensic FACT (genericised, 2026-07-22): a library lock helper opening its lock fd with `exec 9>>"$lock" 2>/dev/null` permanently redirected the INTERACTIVE shell's stderr to `/dev/null` — one provider launch left the operator's terminal unable to print any error, from any command, for the life of the shell — while the

paired unlock helper had used the correct brace form ({ exec 9>&-; } 2>/dev/null) all along; the asymmetry was the defect. Marker-probe proof: pre-fix, a marker echoed to stderr after the call never appeared; post-fix it does. The mandate: in ANY sourced/library shell function, an `exec` whose purpose is only to open/close/move file descriptors MUST brace-scope its side redirections — `{ exec 9>>"$file"; } 2>/dev/null` — so error-suppression applies to the fd-operation alone, never to the enclosing shell; a command-less `exec` combining fd-manipulation operands with an additional unbraced redirection in sourced code is a REVIEW DEFECT (§11.4.142) and a greppable pre-build-gate class. Recommended gate `CM-SHELL-EXEC-REDIRECTION-SCOPED` (scan in-scope sourced/library shell sources for command-less `exec` lines that combine fd-open/close operands with additional unbraced redirections → FAIL citing file:line; brace-scoped forms and `exec` WITH a command are exempt — the §11.4.201(1) false-positive guard) + paired §1.1 mutation (unwrap one brace-scoped fd-open back to the bare form → the gate FAILs; gate-code = separate work item). The instrument-side twin of this construct — the SAME `exec ... 2>/dev/null blinding bash -x` tracing of the code under test — is catalogued in `docs/guides/shell_instrument_footguns.md` per §11.4.201(12).

Captured-evidence enforcement. Pre-build gate `CM-SCRIPT-TARGET-SHELL-PARSEABLE` walks every `.sh` file under the in-scope directories, runs `sh -n` on each with a 30-second per-file timeout, and FAILs if any script fails to parse. The gate's diagnostic output cites the failing file + the `sh -n` error message so the fix location is unambiguous.

Propagation gate. `CM-COVENANT-114-67-PROPAGATION` verifies the §11.4.67 anchor literal is present across the 44-file consumer fleet (parent CLAUDE.md / AGENTS.md + Containers + 10 owned atmosphere submodules + nested SmartTube SharedModules / MediaServiceCore / MSC-SharedModules + 7 HelixQA submodules). Constitution submodule files (`constitution/{Constitution,CLAUDE,AGENTS}.md`) are canonical authority — they carry §11.4.67 by definition.

Paired meta-test mutations. (a) Inject a bash-only construct (`exec > >(/bin/cat) )` outside `eval` into a fresh location inside a test script and assert `CM-SCRIPT-TARGET-SHELL-PARSEABLE` FAILs. (b) Strip `11.4.67` literal from one consumer file and assert `CM-COVENANT-114-67-PROPAGATION` FAILs.

Composes with:

- §11.4.1 (FAIL-bluffs equally forbidden — script-bug FAILs that mask product defects are the exact failure mode §11.4.67 mechanically closes)
- §11.4.4 (test-interrupt-on-discovery — a parse FAIL at cycle start IS the discovery event; cycle stops, fix lands, retest)
- §11.4.6 (no-guessing — `sh -n` produces fact, not opinion)
- §11.4.50 (deterministic consistency — a script that parses on one shell and not another fails the N-iteration identical-result requirement)
- §11.4.51 (live-ADB-first — shell-script fixes are LIVE_ADB_TESTABLE: push the fix, re-run on device, capture evidence before commit)

Mandatory protections (no escape hatch):

- **No `--skip-parseability-check` flag.** A script that does not parse under the target shell is broken regardless of intent.
- **No "this script only runs on bash" exception.** If a script may be invoked via `sh script.sh` (and almost all of them can — Android's default `sh` is mksh), it MUST parse under mksh.
- **No "the runtime guard catches it" exception.** Runtime guards inside a mksh-parsed script cannot prevent mksh parse-time rejection. Use `eval -deferral` or shebang-controlled invocation.
- **Fix at source, not in callsites.** A broken script is fixed in the script itself (per §11.4.1), not by individual callsites defensively detecting parse failures.

Pre-build gate `CM-SCRIPT-TARGET-SHELL-PARSEABLE` + paired mutation.

Propagation gate `CM-COVENANT-114-67-PROPAGATION` + paired mutation.

Canonical authority: constitution submodule `Constitution.md` §11.4.67.

Non-compliance is a release blocker regardless of context.

2026 05 20

§ . . — Positive sink-side /
downstream evidence mandate (User mandate,
— —)

Forensic anchor — direct user mandate (verbatim, 2026-05-20):

"We still do not hear any audio played from D3 device! Arvus Web Dashboard when we play music from D3 shows nothing for Codec In Use! This MUST BE investigated and fixed! How come we passed the tests with Arvus validation? What were values for the Codec In Use field? Empty means nothing! This is not working! It MUST BE FIXED, TESTED AND VERIFIED WITH FULL AUTOMATION TESTING ASAP!!!"

§11.4.13 already mandates consuming the Arvus sink-side report when available. §11.4.68 extends this with a stricter contract: **a test that asserts audio or video routing PASS MUST capture and verify positive sink-side or downstream evidence — never config-only, never metadata-only, never PCM-open-state-only**. The §11.4.13 path described what SHOULD be captured; §11.4.68 fixes the failure mode where tests "PASSed" because the helpers silently SKIPPed the very evidence the assertion depended on.

Positive sink-side / downstream evidence — closed enumeration (at least one MUST be captured for every audio/video routing PASS; none alone is sufficient if the canonical sink/downstream is reachable):

1. **Sink-side codec-state** (Arvus / Sonos / Yamaha / Denon / Marantz / WiSA REST API) showing a **non-empty Codec-In-Use** value matching the test's expected codec regex during the playback window. Empty / `<unreachable>` / `<N.E.>` / `<None>` placeholders are NOT positive evidence.
2. **PCM frames-written delta** from `/proc/asound/cardN/pcmMp/sub0/ status hw_ptr` — the delta over the playback window MUST be strictly positive AND consistent with the expected sample-rate × channel-count × duration. Zero or negative delta is FAIL.
3. **ALSA ELD / EDID-Like-Data** from `/proc/asound/cardN/eld#X.Y` demonstrating the sink negotiated the expected channel count + format. Empty ELD is NOT evidence the sink received audio — it only proves the cable / EDID-channel is up.
4. **ffprobe-on-captured-mp4** for video routing — non-zero frame count, expected codec, expected resolution, expected fps. 0-byte capture (Bug #24 pattern) is FAIL.
5. **Recording-analyzer event match** per §11.4.2 / §11.4.5 timeline — at least one matched event in the analyzer findings.

6. **Tinycap RMS amplitude** for ALSA-loopback audio recordings — non-trivial RMS in the captured WAV (≥ -60 dBFS for line-level playback).

Failure-mode taxonomy — what §11.4.68 specifically forbids:

- `arvus_probe_present` returns 1 → test reports SKIP → suite reports all-green. **FORBIDDEN.** When the sink-side evidence is REQUIRED by the assertion, missing sink is `OPERATOR-BLOCKED` (an explicit release-blocker, NOT a SKIP, NOT a PASS).
- `arvus_probe_codec_state` returns empty string → test logs "codec_state: " (blank) → assertion proceeds on HAL-side metadata. **FORBIDDEN.** Empty codec-state IS a defect signal — either the audio is not arriving at the sink (genuine product defect per current User mandate) OR the sink is unreachable (operator-blocked). Either way, the test cannot PASS.
- Test reads HAL `dumpsys media.audio_flinger` output-thread device field showing `0x400` (`AUDIO_DEVICE_OUT_HDMI`) → reports "HDMI routing PASS". **FORBIDDEN.** That field reflects what the audio POLICY chose, NOT what the audio HAL physically opened. The HDMI ALSA card may still be closed (PCM open failure swallowed by the HAL's `ALOGW` fallback) — silent silence at the user's ear. Sink-side codec-state or PCM `hw_ptr` delta is the ONLY proof audio actually arrived.
- Test asserts `/vendor/etc/audio_policy_configuration.xml` contains the `deep_buffer` route to `HDMI Out`. **FORBIDDEN as the only evidence.** Config-only PASS is the canonical §11.4 PASS-bluff — config can be perfect AND audio still inaudible (e.g., the HAL may still misroute). MUST be paired with at least one runtime positive evidence from the enumeration above.

Mandatory protections (all four):

1. **Library contract.** Sink-side helper libraries (`arvus_probe.sh` , `sonos_probe.sh` , etc.) MUST expose a `*_require_reachable` function that returns exit code 2 (`OPERATOR-BLOCKED`) when the sink cannot be probed. Tests asserting sink-side codec-state MUST call this function at probe entry. Silent skip via `*_probe_present` is permitted ONLY for topology-dispatch fan-out logic that already has an alternative on-SoC + downstream-amplitude probe path.

2. **Exit-code propagation.** `arvus_require_reachable` / equivalent helpers return 2; the test harness MUST propagate 2 to the suite summary (counted separately from FAIL/SKIP) as `OPERATOR-BLOCKED` . `test_all_fixes.sh` orchestrator MUST exit non-zero when any test reported `OPERATOR-BLOCKED` AND release tags MUST NOT be cut while any `OPERATOR-BLOCKED` test exists.
3. **Test rewrite.** Every test currently asserting audio output routing PASS via metadata-only checks (e.g., the historical `test_audio_output_routing.sh` config-only invariants) MUST be rewritten to capture at least one positive sink/downstream evidence per the enumeration above. The §11.4 PASS-bluff pattern `tinymix says route is on → PASS` is the failure mode this anchor closes.
4. **Anti-stickiness post-stop.** After the test stops the playback, it MUST re-probe the sink/downstream and assert the codec-state transitioned to N.E. / N/A / 0-frames-delta. A persistent codec- state across test boundaries indicates orphan playback (§11.4.14 violation) AND would falsely PASS the next test.

Pre-build gate `CM-COVENANT-114-68-PROPAGATION` enforces this anchor literal across canonical files + per-consumer propagation. Paired mutation strips the anchor literal → gate FAILs. Pre-build gate `CM-POSITIVE-SINK-EVIDENCE-PER-AUDIO-TEST` walks every audio-routing on-device test asserting PASS and verifies it cites at least one positive evidence call (one of `arvus_require_reachable` / `arvus_assert_codec_format` / `pcm_hw_ptr_delta_positive` / `tinycap_rms_above_floor` / `ffprobe_frames_nonzero`). Paired mutation strips one such call → gate FAILs.

Composes with §11.4.2 (recorded-evidence — sink-side is the captured-evidence channel), §11.4.5 (audio + video quality analysis comprehensiveness — codec / channel-count / RMS), §11.4.13 (sink-side captured-evidence mandate — §11.4.68 closes its silent-skip gap), §11.4.14 (test cleanup — anti-stickiness post-stop), §11.4.46 (recent-work validation — sink-side evidence required at validation entry), §11.4.49 (dual-approach testing — both UI + Intent variants need sink-side evidence), §11.4.50 (deterministic consistency — sink-side evidence must be consistent across N iterations), §11.4.52 (autonomous-validation — sink probe IS the canonical autonomous path for HDMI audio).

No escape hatch — no `--skip-sink-evidence` , `--allow-empty-codec` , `--sink-unreachable-is-pass` , `--metadata-only-suffices` flag exists anywhere. The discipline exists because the operator/end-user forensic on 2026-05-20 demonstrated the exact failure this anchor forbids: tests reported audio-routing PASS while the user heard nothing and Arvus Codec-In-Use was empty.

Canonical authority: constitution submodule `Constitution.md` §11.4.68.

Non-compliance is a release blocker regardless of context.

2026 05 20

§ . . — Universal Sink-Side Positive-Evidence Taxonomy + Mechanical Enforcement (User mandate, — —)

Forensic anchor — direct user mandate (verbatim, 2026-05-20):

"THIS MUST HAPPEN NEVER AGAIN!!! We MUST HAVE this all working! Not just for audio but for every single piece of the System!!! Proper full automation when executed with success MUST MEAN that manual testing will be as much positive at least regarding the success results! Dive deep into the root causes and how this had happened! The root causes of such omission MUST BE TRACKED, and proper solution applied! Solution MUST BE universal, generic that solves working flows for all System components and for all future and all existing projects! Make sure everything is added we change as mandatory rules we must follow and critical constraints into ours root (constitution Submodule) `Constitution.md`, `CLAUDE.md` and `AGENTS.md`!!! Everything we do MUST BE validated and verified with rock-solid proofs and anti-bluff policy enforcement and fulfillment! THIS IS MOST CRITICAL POINT WE HAVE NOW with all audio issues!"

Escalation from the same operator session (2026-05-20, hours earlier):

"We still do not hear any audio played from D3 device! Arvus Web Dashboard when we play music from D3 shows nothing for Codec In Use! ... How come we passed the tests with Arvus validation? What were values for the Codec In Use field? Empty means nothing! This is not working!"

Forensic incident — 2026-05-19 → 20 D3 audio-routing PASS-bluff generalised.

§11.4.68 closed the audio-specific failure path where Arvus reported an empty Codec-In-Use field and the test PASSEd anyway. The same shape of bug had previously shown up intermittently across multiple subsystems (display routing on D1 secondary HDMI, BT A2DP codec advertisements, WiFi connection quality, video playback frame-count, touch-input event delivery). The audio incident was the canonical trigger because the operator hears silence at the soundbar even while the cycle reports green. §11.4.69 is the universal generalisation — closing the bluff class for every user-visible feature, not only audio.

Root-cause analysis — why existing anchors §11.4.2 / §11.4.5 / §11.4.13 / §11.4.27 / §11.4.52 / §11.4.68 did NOT mechanically catch the failure across the System:

1. **Aspirational text without a pre-build gate.** §11.4.2 (recorded-evidence) required captured visual/audio evidence per PASS but had no gate walking the test corpus to enforce a per-PASS evidence file. Tests called `ab_pass` "description" and the evidence requirement was honour-system.
2. **SKIP-fail-open hiding missing evidence.** §11.4.13 (Arvus sink-side) had a silent-skip escape when the sink returned an empty / unreachable response — the empty response was the very defect signal but it was classified as "sink unreachable, SKIP the assertion." §11.4.68 closed this for audio; §11.4.69 closes it universally across every feature class with a downstream sink.
3. **PASS-by-default vs FAIL-by-default helper contract.** The bare `ab_pass` helper accepted a free-text description and incremented the PASS counter — no evidence path required, no shape enforced. The default failure mode was therefore PASS- bluff. The correct default is FAIL-by-default — a PASS must prove itself with a captured-evidence artefact path.

4. **Missing closed-set taxonomy.** Individual test authors made one-off decisions about what evidence shape was acceptable. Without a single canonical table mapping every user-visible feature class to its required sink-side probe, aggregate enforcement across the System was mechanically impossible.

5. **Helper functions that mapped a multi-step probe to one bool.**

`arvus_probe_present` collapsed reachable-but-empty into the same value as unreachable. The §11.4.68 fix added `arvus_require_reachable` for audio; §11.4.69 codifies the equivalent for every sink (Sonos, Yamaha, Denon, video display analyzer, WiFi sink, BT peer, etc.).

6. **No per-feature evidence ledger.** §11.4.52 supplied an autonomous-validation classification (`AUTONOMOUS_VERIFIED` etc.) but did not require each test to cite a captured-evidence artefact path matching the §11.4.69 taxonomy.

The §11.4.69 mandate — universal, generic, applies to every System component, every owned project, every consumer subscribing via the constitution submodule:

Element 1: Closed-set sink-side / downstream evidence taxonomy. Every user-visible feature class in every consuming project MUST map to exactly one entry in the following table. The taxonomy is the canonical authority on the required evidence shape per feature class. Tests annotate themselves with `# §11.4.69` `FEATURE: <class>` so the pre-build gate can match expected evidence shape.

Feature class	Required probe	Required evidence shape
audio_output	sink-side codec/channel REST probe (Arvus / Sonos / Yamaha / Denon / Marantz / WiSA) + on-device tinycap capture + ffprobe channel assertion	non-empty codec_in_use matching expected codec AND captured WAV with RMS > -60 dBFS AND ffprobe channels matches claim
audio_input	tinycap capture from intended device + ffprobe RMS	captured WAV with RMS > -60 dBFS AND audio_devices_for_attr shows the claimed source device
video_display	screenrecord on intended display + ffprobe -count_frames + analyzer event-match	non-zero mp4 size AND nb_read_frames > 0 AND analyzer reports a matched event for the claimed display
net-work_throughput	iperf3 / curl --speed / dd over network	throughput ≥ per-link-type floor AND TCP buffer / congestion-control matches claim
net-work_connectivity	ping + DNS resolve + TCP connect probe	captured ping reply AND DNS answer AND TCP handshake to known port
bluetooth_a2dp	codec query + sink peer state + A2DP TX queue depth	dumpsys bluetooth_manager reports CONNECTED+A2DP_PLAYING AND codec field non-empty AND queue depth < overflow threshold
bluetooth_pair	dumpsys bluetooth_manager + Bluetooth-Device.getBondState()	bond state = BOND_BONDED AND device MAC matches expected

Feature class	Required probe	Required evidence shape
touch_input	getevent -lt capture during scripted input tap	≥ 1 EV_ABS ABS_MT_* event per scripted tap AND dumphsys input shows the touch device active
sensor	dumphsys sensorservice + getevent on sensor device	event count > 0 between BEFORE and AFTER snapshots for the claimed sensor
gpu_render	SurfaceFlinger frame-rate dump + dumphsys gfxinfo <pkg>	Janky frames < threshold AND Total frames rendered > 0 during test window
storage_read	dd if=<path> of=/dev/null + /proc/diskstats delta	throughput $\geq$ floor AND read_sectors delta confirms IO landed on intended device
storage_write	dd of=<path> + sync + /proc/diskstats delta	throughput $\geq$ floor AND write_sectors delta confirms IO landed on intended device
mediaco-dec_decode	media.metrics dump + dumphsys media.codec	decoder lifecycle events present (configure , start , ≥ 1 output buffer) AND codec name matches claim
mediaco-dec_encode	media.metrics dump + non-zero output file	encoder lifecycle events present AND output file non-zero AND fprobe confirms codec matches claim
miracast / cast	sink-side WFD / Cast announce + session state	sink dashboard reports active session AND session ID present in dumphsys media_session

Feature class	Required probe	Required evidence shape
boot_service	getprop init.svc.<name> + log capture	service state = running AND log shows expected boot-time output literal
package_install	pm path <pkg> + signature match	non-empty pm path AND APK signature matches ex- pected platform key
permis- sion_grant	dumpsys package <pkg> permission state	permission state = granted: true for the claimed permission
wifi_link	iw dev wlan0 link + RSSI + frequency	captured BSSID AND RSSI within range AND frequency matches claim
wifi_throughput	iperf3 to LAN server	throughput ≥ floor for the ne- gotiated PHY rate
ethernet_link	ip link show + carrier state + throughput probe	carrier=1 AND speed matches claim AND through- put probe succeeds
dis- play_topology	DRM connector dump + dis- play ID enumeration	connector status connected for claimed connector AND display ID present in dumpsys Sur- faceFlinger
drm_playback	Widevine session + decoded- frame count	active DRM session AND decoded-frame count > 0
subtitle_render	analyzer + Tesseract OCR on captured frames	OCR matches expected subtitle text on the claimed display

The taxonomy is **open to additions** via constitution-submodule commits + propagation. Removal of a feature class is forbidden without a tracked work item explaining the supersession. Consumer projects MAY extend the taxonomy with project-specific feature classes (e.g., a consuming project's build-pipeline classes) but MUST NOT contract it.

Element 2: New helper contracts (additive during grace period; mandatory after grace-period end-date 2026-06-19).

The project anti-bluff helper library (`anti_bluff.sh` in a consuming project; functionally equivalent helper in every consuming project) MUST provide three new helpers:

1. **`ab_pass_with_evidence <description> <evidence_path>`** — the new canonical PASS helper:
 - REQUIRES a non-empty `evidence_path` argument
 - Verifies the path exists AND is non-empty (`[ -s "$path" ]`)
 - If the path is missing or empty, the helper EXITS the test as FAILED with a diagnostic citing the missing evidence path
 - On success, increments the PASS counter AND emits a result line `PASS: <description> [evidence: <path>]`
 - Composes with `ab_send_action` — a single test issues many `ab_send_action` calls and one terminal `ab_pass_with_evidence`
2. **`ab_skip_with_reason <description> <reason_code>`** — the new canonical SKIP helper:
 - REQUIRES `reason_code` from a CLOSED set: `geo_restricted` , `operator_attended` , `hardware_not_present` , `topology_unsupported` , `network_unreachable_external` , `feature_disabled_by_config` , `artifact_not_yet_built`
 - REJECTs ad-hoc reasons — any string not in the closed set fails the helper at call time
 - `artifact_not_yet_built` (added 2026-07-17, research-derived) names the NOT-YET-RUNNABLE class: the check is CORRECT and the topology IS present, but its precondition artifact does not exist yet (the fix is not built / not deployed to the target). It is the legal string for the §11.4.135 REQUIRES-REBUILD PENDING class, which previously had NO code in

this closed set and therefore had to either misuse a false code (`hardware_not_present` / `feature_disabled_by_config` — both untrue, and a §11.4.6 misstatement of the real reason) or fail the helper at call time. Reporting a not-yet-runnable check as FAIL is a §11.4.1 FAIL-bluff (it sends teams to fix healthy code); reporting it as PASS is a §11.4 PASS-bluff. Its seam-dependence is §11.4.135's, NOT flat: on an intermediate artifact it is honest and non-blocking; on the release candidate it is a contradiction that BLOCKS (absence of a verdict blocks exactly as a FAIL does). A flat "always blocks" rule is REJECTED per §11.4.135 as itself a §11.4.201 FAIL-bluff. Consumer-migration note (§11.4.164 seam): adding a member to a CLOSED set means every existing `ab_skip_with_reason` implementation REJECTS the new code at call time until updated — update the helper alongside the constitution pull.

- FORBIDS `network_unreachable_external` for any feature class in the §11.4.69 taxonomy that lists a sink-side probe — that escape was the §11.4.13 fail-open path §11.4.68 closed for audio, and §11.4.69 closes universally
- SKIP remains mechanically distinct from PASS in cycle math (per §11.4.1)

3. **Bare `ab_pass` deprecation wrapper.** The legacy `ab_pass` helper MUST emit a WARNING line on every call AND emit a deprecation event to a tracked log file (`/data/local/tmp/ab_pass_deprecations.log`). After the grace period (defined in Element 4) the wrapper MUST FAIL the test outright, forcing all callers to migrate.

Element 3: Mechanical enforcement via three pre-build gates + three paired meta-test mutations (per §1.1).

1. **CM-SINK-EVIDENCE-PER-FEATURE** — walks every `.sh` test in `device/rockchip/rk3588/tests/` (and equivalent project paths), parses for `# §11.4.69 FEATURE: <class>` annotation comments, verifies the test invokes the corresponding sink-side probe from the taxonomy AND uses `ab_pass_with_evidence` (or `ab_skip_with_reason` with a permitted reason). Pre-grace: tests lacking the annotation WARN. Post-grace: they FAIL.
2. **CM-NO-FAIL-OPEN-SKIP** — audits sink-side probe helpers (`arvus_probe.sh` , `sonos_probe.sh` , equivalents) for any code path that converts an empty / unreachable sink response into a PASS-counting SKIP for a feature class with a

sink-side probe in the taxonomy. The gate FAILs if any such path exists.
Remediation: escalate to FAIL (the empty response IS the defect) or to OPERATOR-BLOCKED per §11.4.68.

3. **CM-AB-PASS-WITH-EVIDENCE-EVERYWHERE** — pre-grace WARNs on every bare `ab_pass` call in in-scope tests; post-grace (2026-06-19) FAILs the gate. Tests that legitimately have no evidence-capable PASS path MUST migrate to `ab_skip_with_reason operator_attended` and add a tracked work item per §11.4.52's `OPERATOR_ATTENDED_ONLY` classification.

Paired mutations (per §1.1):

- **Mutation 1:** strip the literal `ab_pass_with_evidence` from the project anti-bluff library → assert `CM-AB-PASS-WITH-EVIDENCE-EVERYWHERE` FAILs.
- **Mutation 2:** inject a fresh test that calls bare `ab_pass` without the `# §11.4.69 FEATURE: annotation` → assert `CM-SINK-EVIDENCE-PER-FEATURE` FAILs post-grace, WARNs pre-grace.
- **Mutation 3:** inject a fresh test that calls `ab_skip` (no `_with_reason` suffix) OR `ab_skip_with_reason network_unreachable_external` for an `audio_output` feature → assert `CM-NO-FAIL-OPEN-SKIP` FAILs.

Element 4: Grace period mechanics — additive, not destructive.

- **Grace period:** 30 days from anchor land date — 2026-05-20 through 2026-06-19 inclusive. The end date is hardcoded in the helper library and in this anchor; no environment variable, no command-line flag, no operator override can extend or shorten it.
- **Pre-grace behaviour:** new helpers exist and are usable; the three gates emit WARN for legacy patterns; cycle math counts WARNs separately from PASS / FAIL / SKIP.
- **Post-grace behaviour:** WARN promotion to FAIL is automatic via a date check `( [ "$(date -u +%Y%m%d)" -ge 20260619 ] )` inside each gate's enforcement logic.

Element 5: Inheritance via HelixConstitution. This anchor lives in the constitution submodule and applies to every consuming project that subscribes — including Catalogizer and every future consumer. Each consumer MUST implement its own

equivalent helper library + pre-build gates; the taxonomy in Element 1 is the canonical authority across all consumers. Consumer-specific extensions are allowed only as additions to the taxonomy.

Composes with:

- §11.4.1 (FAIL-bluffs equally forbidden — `ab_skip_with_reason` closed-set prevents script-bug FAILs hiding real defects)
- §11.4.2 (recorded-evidence requirement — Element 2 helper contracts make captured evidence mechanically required per PASS)
- §11.4.5 (audio + video quality 5-layer — §11.4.69 enforces that the layers are exercised before PASS via the taxonomy probe)
- §11.4.6 (no-guessing — sink-side evidence eliminates "we think audio is routing" guessing; the codec field reads concretely)
- §11.4.13 (Arvus sink-side evidence — §11.4.69 mechanically closes the SKIP-fail-open escape §11.4.13 left open)
- §11.4.27 (no-fakes-beyond-unit — §11.4.69 adds the runtime hook §11.4.27 lacked)
- §11.4.50 (deterministic consistency — sink-side evidence is what `ab_run_n_times` hashes for divergence detection)
- §11.4.52 (autonomous-validation — §11.4.69 supplies the evidence-shape contract per feature class; §11.4.52 supplies the classification ledger)
- §11.4.68 (positive sink-side / downstream evidence for audio + video — §11.4.69 is the universal generalisation across all feature classes)

No escape hatch:

- No `--skip-evidence` flag.
- No `--config-only-pass` flag (configuration-only PASS is the exact bluff pattern this anchor closes).
- No `--allow-fail-open-skip` flag.
- No `--legacy-ab-pass-permitted` flag post-grace.
- No taxonomy-bypass shortcut — every feature class in scope MUST have its evidence shape enforced.

Propagation gate. `CM-COVENANT-114-69-PROPAGATION` verifies the §11.4.69 anchor literal is present across the ~44-file consumer fleet (parent `CLAUDE.md` / `AGENTS.md` + Containers + 10 owned atmosphere submodules including the smarttube-player nested set + 7 HelixQA submodules). Constitution submodule files (`constitution/{Constitution,CLAUDE,AGENTS}.md`) are canonical authority — they carry §11.4.69 by definition.

Canonical authority: constitution submodule `Constitution.md` §11.4.69.

Non-compliance is a release blocker regardless of context.

§ . . — Subagent-Driven Execution Is The Default (User mandate, — —)

Forensic anchor — direct user mandate (verbatim, 2026-05-20):

"Always do if possible Subagent-driven! Add this into our root (constitution Submodule) `Constitution.md`, `CLAUDE.md` and `AGENTS.md`. This should be the default choice ALWAYS!"

When executing implementation plans authored via `superpowers:writing-plans` (or any equivalent task-decomposed execution flow), the **default execution model is subagent-driven** per `superpowers:subagent-driven-development`. Inline execution via `superpowers:executing-plans` is permitted ONLY when (a) the task is trivial AND fits in a single sub-300-line edit, OR (b) the operator explicitly requests inline execution at brainstorm-handoff time per the `superpowers:writing-plans` skill's explicit-execution-mode prompt.

Why subagent-driven is the default:

1. **Isolated context per task** — each subagent gets a fresh context window scoped to its task, avoiding the conductor's context bloat that degrades reasoning quality after ~100 k tokens.
2. **Two-stage review naturally enforced** — fresh subagent does the work, conductor reviews, optional second subagent verifies. This composes with §11.4.4 four-layer coverage + §11.4.43 TDD-fix- discipline.

3. **Parallel-PWU compatible** (§11.4.58) — subagents are the unit of parallelism in the PWU pipeline; making them the default aligns the writing-plans handoff with the parallel-development methodology.
4. **Anti-bluff seam** (§11.4) — the conductor's review of a subagent's output is structurally separated from the work itself, eliminating self-review blind spots.
5. **Survives operator absence** — subagents resume from their on-disk plan + spec inputs; the conductor session can be lost without losing in-flight task state.

Composition with §11.4.4 (four-layer coverage), §11.4.6 (no- guessing — subagent's captured output IS the evidence), §11.4.42 (iteration discipline), §11.4.43 (TDD-fix), §11.4.50 (deterministic consistency — the subagent's deterministic exit code), §11.4.51 (LIVE_ADB_FIRST — subagents classify rebuild-requirement before commit), §11.4.58 (parallel-development PWU — subagents ARE the parallel work units).

No escape hatch. `--inline-execution-required` , `--no-subagents` , `--monolithic-execution` are NOT permitted flags. Operator may request inline at brainstorm-handoff time per `superpowers:writing-plans` skill; absent that, subagent-driven is default. Skipping subagent-driven for non-trivial work without recorded operator authorisation is itself a §11.4 PASS-bluff (the conductor self-reviews, the anti-bluff seam disappears, and the forensic class of incident this Constitution catalogues becomes mechanically possible again).

Captured-evidence enforcement. Pre-build gate `CM-COVENANT-114-70-SUBAGENT-DEFAULT-PROPAGATION` enforces this anchor literal in every CLAUDE.md / AGENTS.md across parent + 10 owned submodules + nested submodules + HelixQA dependencies (the ~44-file consumer fleet). Paired §1.1 meta-test mutation strips the anchor literal → gate FAILs.

Canonical authority: constitution submodule `Constitution.md` §11.4.70. Default applies to every `superpowers:writing-plans` skill terminal handoff and every multi-task execution flow with independent subtasks.

Non-compliance is a release blocker regardless of context.

§ . . — Pre-Push Fetch + Investigate + Integrate Mandate (User mandate, — —)

Forensic anchor — direct user mandate (verbatim, 2026-05-20):

"before pushing changes to any upstream for any repository - main repo or Submodule, we MUST fetch and pull all changes. Once these are obtained WE MUST investigate what is different compared to head position we were on last time before fetching and pulling new changes! We MUST understand what is done and for what purpose, easpecially how that does affect our project and our System in general! Any mandatory changes or improvements required by fresh changes we just have brought in MUST BE incorporated, covered with all supported types of the tests which will produce as a result of its success execution REAL PROOFS of working for all componetns and functionalities covered and work fully in anti-bluff manner!"

This anchor is the **everyday-push variant** of §11.4.41 (Pre-Force-Push Merge-First). §11.4.41 governs the destructive force-push case; this anchor governs EVERY push including ordinary fast-forwards. The complementary discipline closes the regression-via-stale-baseline gap: a project committing against a HEAD that's behind any upstream's main risks pushing changes that conflict with, regress, or fail to incorporate accepted work landed in parallel by other consumers (Catalogizer and other consuming projects on the constitution submodule; sibling projects on other shared submodules).

The mandatory 5-step pre-push cycle (every push, every repository — main + every submodule):

1. **Fetch all remotes** — `git fetch --all --prune --tags` (or per-remote loop if `--all` is unavailable). Capture stdout for the audit trail.
2. **Pull all upstream branches** — `git pull --no-rebase <remote> <branch>` for each configured remote whose tip differs from the local branch. Resolve any merge conflict per §11.4.41 step 2 (rebase / merge / cherry-pick per consumer judgment, NOT auto- `--ours` or `--theirs`).

3. **Investigate the diff vs OUR previous HEAD** — `git log --oneline <previous-head>..HEAD + git diff --stat <previous-head>..HEAD +` read EVERY commit's body. For each foreign commit understand: (a) what changed, (b) why (forensic anchor + cited user mandate or §-anchor), (c) how it affects OUR project + OUR System (does it touch surfaces we ship? does it introduce constraints we now MUST honor? does it deprecate or supersede patterns we use?).
4. **Integrate mandatory changes + cover with full anti-bluff test coverage** — if the pulled-in work mandates anything (new §-anchor, new gate, new propagation requirement), implement it per §11.4.4(b) four-layer coverage (pre-build gate + post-build inspection + on-device test + HelixQA Challenge) + §11.4.43 TDD-fix RED-test-first discipline. Every PASS MUST carry §11.4.5 captured-evidence (REAL PROOFS — not metadata-only).
5. **Then push** — only after Steps 1–4 land. Push to every configured remote in the per-repo cascade order (canonical first, then mirrors). Verify the push landed with `git ls-remote <remote> <branch> post-push`.

Composition with §11.4.26 (constitution-submodule update pipeline — per-submodule specialisation) + §11.4.32 (post-pull validation) + §11.4.37 (fetch-before-edit) + §11.4.40 (full-suite retest before tag) + §11.4.41 (pre-force-push merge-first — the force-push case of this anchor) + §11.4.42 (iteration discipline) + §11.4.43 (TDD-fix-discipline) + §11.4.4(b) (four-layer coverage) + §11.4.5 (audio + video quality analysis comprehensiveness) + §11.4.6 (no-guessing — every integration decision cites the foreign commit by SHA, not "we think it does X").

No escape hatch — there is no `--skip-fetch`, `--no-investigate`, `--fast-push`, `--trust-upstream` flag. The discipline exists because:

- Pushing-without-fetching produces silent stale-baseline regressions (consumer A's PR diverges from consumer B's PR; neither saw the other; both merge; resulting tree corrupts).
- Pulling-without-investigating destroys the foreign work's traceability (we now carry someone else's code without understanding it; future debugging crosses the §11.4.6 no-guessing line).
- Integrating-without-testing reproduces the §11.4 PASS-bluff pattern (we adopt a change but don't prove it works in OUR system).

- Pushing-after-integration-without-evidence reproduces §11.4 itself (untested integration = release-blocker by definition).

Per-repository scope: this anchor applies to (a) the parent consuming project's repo, (b) the constitution submodule (across all 6 remotes), (c) every owned submodule (presenter, vlc-player, nova-player, mpv-player, gramophone-player, rhythm-player, strep-player, smarttube-player, torrservice, lampa), (d) every nested submodule (smarttube-player/SharedModules / smarttube-player/MediaServiceCore / smarttube-player/MediaServiceCore/SharedModules), (e) every HelixQA dependency (Challenges, Containers, DocProcessor, LLMOrchestrator, LLMProvider, VisionEngine, HelixQA).

Audit-trail requirement: every push event MUST be reconstructable post-hoc from `docs/changelogs/<tag>.md` + the per-repo `git log` evidence. Operators auditing a release MUST be able to: (1) see what each repo's `HEAD..@{u}` was BEFORE this push, (2) see what foreign work was pulled in, (3) see what new tests + captured-evidence were generated by the integration. Anchor enforcement may be partially automated by extending `scripts/commit_all.sh` to refuse pushes without an audit-trail log entry.

Pre-build gate `CM-COVENANT-114-71-PROPAGATION` enforces this anchor literal across the consumer fleet (parent + 10 owned submodules + nested + HelixQA). Paired mutation strips the anchor literal → gate FAILs. Paired meta-test also asserts the audit-trail-log file produced by a recent push exists + lists the foreign-commit SHAs that were pulled in.

Canonical authority: constitution submodule `Constitution.md` §11.4.71.

Non-compliance is a release blocker ²⁰²⁶⁻⁰⁵⁻²⁰ regardless of context.

§ . . – Audio Top-Priority Mandate
(User mandate, – –)

Forensic anchor — direct user mandate (verbatim, 2026-05-20):

"Make sure all fixes for audio are always top priority in main working stream!"

The conductor (main working stream — whether a Claude Code session, an AI agent, or a human operator) MUST treat audio fixes as the highest-priority class on the serial dispatch queue. Audio fixes include but are not limited to §EU HDMI rejection, §EM multichannel HDMI, §ET Arvus integration, §EV/§EW/§EX D3 audio defects, HiFi (Fix #74 pcm_config_deep, #76 HDMI/SPDIF resilience), AC3 (Fix #106 FFmpeg Codec2), ES8388 (Fix #103/§104/#105 mutual-output + DAPM), multichannel LPCM (Fix #112), and every future audio-stack improvement.

Composition with §11.4.42 (iteration-discipline mandate) + §11.4.58 (parallel-development methodology). Per §11.4.42, the conductor selects work items in priority order; this anchor pins audio at the apex of that order. Per §11.4.58 PWU pipeline, parallel BACKGROUND subagents (research, refactors, infrastructure documentation) MAY run concurrently with audio work — they do NOT preempt audio on the main-stream serial dispatch queue.

Operative rule: any time the conductor faces a choice between dispatching an audio task vs a non-audio task on the SAME serial resource (most commonly the pre_build_verification.sh + meta_test_false_positive_proof.sh edit-and-commit critical path), the audio task wins.

No escape hatch — there is no "but this non-audio task is faster" or "but this research is more interesting" override. The discipline exists because audio-stack regressions are user-perceptible and high-impact (D3 silent post-flash is a release blocker), while research and refactors can wait.

Canonical authority: constitution submodule `Constitution.md` §11.4.72.

Non-compliance is a process violation regardless of context.

2026 05 20

§ . . — Main-specification document
versioning + revision discipline (User mandate,
– –)

Forensic anchor — direct user mandate (verbatim, 2026-05-20):

"Make sure everything we add now in previous and upcoming requests IS ALWAYS applied to the main specification — if we have one. Since all these are not major changes we could increase Specification version per change for secondary version instead of the primary. Primary version MUST BE increased for much bigger levels of changes! Add this into root (constitution Submodule) Constitution.md, CLAUDE.md and AGENTS.md as mandatory rule / constraint applicable ONLY IF we have something like the main specification document or we do recognize something like the main specification document. Document MUST BE updated ALWAYS to follow the versioning rules we are applying here + revision and other properties we have!"

Scope condition. This clause applies **only when a project recognises a main specification document** — i.e. a project-level Markdown file (or set of files) that constitutes the canonical "what this system does" reference, typically at `docs/specs/**/specification*.md` or a comparable path. Projects with no main specification (small libraries, internal tooling) are exempt. Projects that operate per-feature spec files (e.g. ADRs only) are exempt unless they elect to adopt the main-spec pattern.

The mandate. When a project DOES have a main specification document:

1. **Every additive operator requirement, refinement, or accepted recommendation MUST be applied to the spec** before, or as part of, the work that implements it. The spec is the source of truth; downstream code and tracking docs reference it.
2. **Spec versioning has two axes** — *primary* and *secondary* (where "secondary" maps to the existing §11.4.61 metadata-table `Revision` row):
 - **Primary (V1 / V2 / V3 / ...)** bumps for major rewrites — substantial architecture changes, foundational scope shifts, deprecating large sections, replacing the technology stack. Triggered explicitly by operator decision; old primary versions move to `archive/` per §11.4.61.

- **Secondary (Revision)** bumps for every other change — section additions, mandate landings, structural reorganisation, polish, type-contract refinements. The metadata table's `Revision` integer is the secondary version (matches §11.4.44 monotonic-integer rule).

3. **The metadata table on the spec MUST stay current** — `Revision` bumps in lockstep with the change; `Last modified` updates to the change date; `Status summary` describes the bump's content; `Fixed` and `Fixed summary` reference the relevant work IDs.
4. **Cross-document propagation** — when a project has propagated copies of the spec-change rule (Herald uses `CLAUDE.md`, `AGENTS.md`, `HERALD_CONSTITUTION.md` per its §106), those copies MUST also reference the active spec file (`specification.V<primary>.md`) and not a stale archived version.
5. **Archive semantics** — when the primary version bumps, the old `specification.V<n>.md` moves to `<spec-dir>/archive/specification.V<n>.md` and gains `Status: superseded` with a `Continuation` pointer to the new active file. Per the §11.4.65 universal-export mandate the archived `.html` and `.pdf` siblings travel with it.

Anti-bluff captured-evidence gate (planned). `CM-SPEC-VERSION-DISCIPLINE` :

- If the project advertises a main spec (operator marks via `[project].main_spec_path = "docs/specs/.../specification.V<n>.md"` in a canonical config file), the gate asserts:
 - The file at `main_spec_path` exists.
 - Its metadata table's `Revision` row is $\geq$ the highest `Revision` referenced anywhere else in the repo (`Issues.md`, `Fixed.md`, `CLAUDE.md`, `AGENTS.md`).
 - If commits exist that touch the spec since the last `Revision` bump, the gate FAILs (operator forgot to bump).

Paired §1.1 mutation. Backdate the `Last modified` row to a prior date while leaving real-content edits in place → gate FAILs.

Composition. Composes with §11.4.44 (revision header — `Revision` is the secondary version here), §11.4.61 (metadata table + ToC), §11.4.59 (README always-sync — if README cites the spec, the citation MUST point at the current primary version), §11.4.65 (universal export — archived versions stay exported).

Classification: universal (per §11.4.17), applicable conditionally per the scope condition above.

Canonical authority: constitution submodule `Constitution.md` §11.4.73.

^{11 4 74}
Non-compliance (forgetting to bump, or letting the spec drift behind operator-mandated requirements) is a release blocker.

^{2026 05 20}

§ . . — Submodule-catalogue-first
discovery + extend-don't-reimplement (User
mandate, — —)

Forensic anchor — direct user mandate (verbatim, 2026-05-20):

"We MUST ALWAYS check which already developed features / functionalities do exist as a part of our comprehensive Submodules catalogue located in vasic-digital and HelixDevelopment organizations on GitHub and GitLab both! Project MUST BE aware of all its existence so we do not implement same things multiple times if they are already done as some of existing universal, reusable general development purpose Submodules! For any missing features that some Submodules we incorporate may be missing we MUST IMPLEMENT the properly and extend those Submodules further! We do control all of the and we CAN and MUST maintain and extend the regularly! All development cycle rules we have MUST BE applied to them and fully respected!"

The mandate. Before scaffolding ANY new module, package, helper, or utility, the contributor (human or AI agent) MUST:

1. **Survey the canonical Submodule catalogue.** Two organisations are authoritative:
 - `vasic-digital` on GitHub (<https://github.com/vasic-digital>) AND GitLab (<https://gitlab.com/vasic-digital>).

- `HelixDevelopment` on GitHub
(<https://github.com/HelixDevelopment>) AND GitLab (<https://gitlab.com/HelixDevelopment>). Both organisations mirror across all four supported hosts per §2.1 — surveying one mirror MUST yield the full catalogue.

- 2. Inventory existing Submodules.** The canonical repo INDEX is `submodules-catalogue.md` in this submodule (landed 2026-05-20 per user follow-up mandate). It enumerates every owned-by-us repository under `vasic-digital + HelixDevelopment` (142 repos at landing time) with a one-line capability description, categorised so consuming projects can locate "is there already a Submodule for X?" in one glance. The catalogue MUST be re-generated quarterly (or after any non-trivial repo create/rename) per its §3 regeneration recipe.
- 3. Reuse before reimplement.** If a Submodule already provides the functionality (or 80%+ of it), the consuming project MUST add the existing Submodule as a Git submodule, not write the functionality fresh.
- 4. Extend in-place when 80%+ matches but features are missing.** When an existing Submodule is close-but-not-complete, the missing features MUST be added TO THAT SUBMODULE (PR upstream + bump consuming project's submodule pointer) — never as a separate consuming-project helper that duplicates the 80% already in the Submodule.
- 5. Same development-cycle rules apply.** Every Submodule under the two orgs is subject to: §11.4 anti-bluff covenant, §1.1 paired-mutation tests, §11.4.10 credentials handling, §11.4.61 metadata table + ToC, §11.4.65 universal Markdown export, §11.4.73 spec versioning (if the Submodule has its own spec), §2.1 multi-mirror push, §3 propagation order. Maintenance + extension of those Submodules MUST follow these rules without exception.
- 6. Document the survey result.** Each new feature's tracker entry (Issues.md row, ADR, PR description, or equivalent) MUST contain a `Catalogue-Check:` field with one of three values:
 - `Catalogue-Check: reuse <org/repo>@<sha>` — the existing Submodule covers this.
 - `Catalogue-Check: extend <org/repo>@<sha>` — extending an existing Submodule; reference the upstream PR.

- `Catalogue-Check: no-match <date>` — surveyed both orgs on the named date; no existing Submodule covers this functionality; the new code is justified to live in-project. Operators reviewing the PR re-validate this claim.

Anti-bluff captured-evidence gate (planned). `CM-SUBMODULE-CATALOGUE-CHECK` :

- For every new module / package introduced in a non-trivial PR (heuristic: ≥ 200 LOC of new non-test code), the gate scans the PR description and any linked tracker row for a `Catalogue-Check:` line. Absence is a FAIL.
- The gate's paired mutation strips the line from a PR and asserts the gate FAILs.

Composition. Composes with §3 (submodule commits propagate first), §4 (every tag on the main repo MUST be mirrored on every owned submodule), §11.4.36 (mandatory `install_upstreams` on clone/add), §11.4.31 (Submodule-Dependency-Manifest), §11.4.32 (post-pull validation), §11.4.28 (Submodules-As-Equal-Codebase).

Why this matters. Without the catalogue-check, consuming projects re-implement UUIDv7 helpers, RLS-tenant-context wrappers, Pandoc-multi-format exporters, `install_upstreams` shell scripts, Markdown ToC generators, `fnv1a32` hash wrappers, etc. — fragmenting the catalogue and forcing the next project to choose between three subtly-different implementations of the same thing. The mandate makes the catalogue load-bearing.

Classification: universal (per §11.4.17). Applies to every project that consumes this Constitution.

Canonical authority: constitution submodule `Constitution.md` §11.4.74.

Non-compliance is a process violation; severe cases (duplicate implementation landed without catalogue check) are release blockers.

2026 05 20

§ . . — Mechanical Enforcement Without Exception (User mandate, — —)

Forensic anchor — direct user mandate (verbatim, 2026-05-20):

"Why do these violations still happen!? This is a serious problem! We cannot rely on stability nor consistency if we cannot respect our Constitution, mandatory rules and constraints! Is there a way to make this always respected, followed and applied without exception fully and unconditionally!? WE MUST HAVE THIS WORKING FLAWLESSLY!!! Do investigate the root causes of such problems! Once all problems are identified WE MUST apply proper mechanisms for this not to happen NEVER EVER AGAIN!"

§-slot history note. This anchor was originally drafted as §11.4.74 but renumbered to §11.4.75 per the §11.4.41 merge-first mandate after upstream landed a different §11.4.74 (Submodule-catalogue- first discovery) concurrent with this work. Same renumber pattern as Phase 39.DH §11.4.41 collision resolution (commit `8ad51547115` of the parent consuming project's repo).

The §11.4 covenant ("tests pass while feature broken-for-end-user") historically relied on agent and operator vigilance to enforce. Three forensic incidents in 2026-05-19 → 20 demonstrated the failure: subagent `a84d2c86f90fdc297` committed `docs/research/workstation/Configurations.md` without HTML+PDF siblings (§11.4.65 violation) when it stalled at the 600s watchdog before its pandoc-export step; subagent `a6b38eedf4ef796e8` dispatched to remediate ALSO stalled at the same operational seam; the parent CLAUDE.md / AGENTS.md drift after §11.4.66 propagation required a dedicated catch-up commit `f93b25a92eb` . The common failure mode: late-binding enforcement that fires at `pre_build_verification.sh` time — hours-to-days after the violator commit reached every remote.

The mandate. Constitutional invariants MUST be enforced mechanically at FIVE independent layers — bypassing any single layer does not bypass the discipline. The five layers are:

1. **Local pre-commit git hook** — refuses staged `.md` lacking sibling `.html` + `.pdf` (and other staged-only invariants).
2. **commit_all.sh integration** — the canonical project commit script invokes the same checks + auto-runs `sync_all_markdown_exports.sh` to self-repair before the commit.

3. **Local pre-push git hook** — re-runs siblings + propagation gate subset on every commit in the push.
4. **post-commit auto-repair hook** — detects orphan `.md` in the just-committed manifest, auto-generates siblings via pandoc + weasyprint, creates a `chore($11.4.75): auto-export ...` follow-up commit. Idempotent + recursion-guarded.
5. **Local-only equivalent** (Phase 39.GF, User mandate 2026-05-20). Remote CI surfaces (GitHub Actions, GitLab pipelines, Jenkins, CircleCI, etc.) are DISABLED per User mandate 2026-05-20 — the workflow file is preserved at `.github/workflows/constitution-compliance.yml.disabled-local-only` so its contents survive for re-enable but GitHub Actions does NOT execute it (only `.yaml` / `.yml` files under `.github/workflows/` are honoured). Layer 5 enforcement migrated to the LOCAL pre-build- verification + meta-test run that the operator MUST execute before tagging: `bash device/rockchip/rk3588/tests/pre_build_verification.sh`
 - `bash scripts/testing/meta_test_false_positive_proof.sh` . Layers 1-4 remain authoritative; Layer 5 is now the operator's local final- gate ritual. A future re-enable PWU may re-establish remote CI.

Helper contracts (mandatory):

- `scripts/install_git_hooks.sh` — idempotent installer that copies hooks from tracked `scripts/git_hooks/` into `.git/hooks/` . Hooked into `scripts/setup.sh` so fresh clones auto-install. Verifies itself with `--verify` flag.
- `scripts/git_hooks/pre-commit` — Layer 1 implementation.
- `scripts/git_hooks/pre-push` — Layer 3 implementation.
- `scripts/git_hooks/post-commit` — Layer 4 implementation.
- `scripts/git_hooks/commit-msg` — Bypass-rationale enforcement.
- `_constitution_sibling_check` function in `scripts/commit_all.sh` — Layer 2 implementation.

Bypass policy. No layer is perfectly bypass-proof in isolation — that is the design. Bypassing ALL FIVE simultaneously is mechanically impossible because:

- Layers 1 + 3 can be bypassed by `--no-verify`. Layer 5 (now local pre-tag ritual after Phase 39.GF) re-runs the same check before tag creation; tags are NEVER created without it per §11.4.40.
- Layer 2 can be bypassed by direct `git commit`. Layer 1 catches that.
- Layer 4 cannot be bypassed by `--no-verify` (runs after commit lands).
- Layer 5 is now local-only (Phase 39.GF — User mandate 2026-05-20). The operator runs `pre_build_verification.sh` + meta-test before every tag per §11.4.40 full-suite-retest mandate. Skipping the local Layer 5 ritual blocks tag creation.

Legitimate bypasses MUST be audited. When `--no-verify` is detected via touch-file marker `.git/ATMO_LAST_BYPASS_ATTEMPT`, the `commit-msg` hook REQUIRES a `Bypass-rationale: <reason>` footer in the commit message. `docs/audit/bypass_events.md` accumulates the audit trail per §11.4 captured-evidence requirement.

Captured-evidence enforcement. Five pre-build gates with paired §1.1 meta-test mutations:

- `CM-COVENANT-114-75-PROPAGATION` — anchor literal across canonical files.
- `CM-GIT-HOOKS-INSTALL-SCRIPT` — installer present + executable.
- `CM-GIT-HOOKS-SOURCE-DIR` — 4 hook bodies present + executable.
- `CM-COMMIT-ALL-SIBLING-CHECK` — `_constitution_sibling_check` in `commit_all.sh`.
- `CM-CI-WORKFLOW-PRESENT` — CI workflow + ≥3 required jobs.

Composes with §1.1 (paired meta-test mutations), §9 (data safety), §11.4 (end-user-quality covenant — Layer 5 CI enforcement is the §11.4 promise made mechanical), §11.4.1 (FAIL-bluffs forbidden), §11.4.4 (test-interrupt-on-discovery), §11.4.5 (audio + video quality), §11.4.6 (no-guessing), §11.4.41 (pre-force-push merge-first — this very anchor's renumber from §11.4.74 → §11.4.75 was driven by §11.4.41 merge-first discipline), §11.4.43 (TDD-fix), §11.4.51 (live-ADB-first), §11.4.52 (autonomous-validation), §11.4.58 (PWU pipeline), §11.4.65 (universal Markdown export — Layer 1+3+4 are §11.4.65's mechanical seam), §11.4.66 (interactive-clarification), §11.4.67 (target-shell- parseability — hooks parse under

bash AND mksh), §11.4.69 (universal sink-side positive-evidence), §11.4.70 (subagent-driven execution), §11.4.71 (pre-push fetch + integrate — Layer 3 + 5 honour the merge- first pipeline), §11.4.72 (audio top-priority — hooks do not race against in-flight audio commits because they hold no lock themselves), §11.4.73 (main-spec versioning — when spec exists, hooks include it), §11.4.74 (submodule-catalogue-first — hooks can be added to the canonical catalogue for cross-project reuse).

No escape hatch. No `--skip-hooks`, `--bypass-enforcement`, `--allow-orphan-md`, `--ci-not-applicable`, `--mechanical-enforcement-not-needed` flag exists. The `--no-verify` route IS the deliberate audit-trail bypass; §11.4.75 makes the audit trail mechanical via the `commit-msg` footer requirement.

The mandate exists because the User mandate of 2026-05-20 is unambiguous: violations MUST NEVER EVER AGAIN happen. Reliance on vigilance has demonstrably failed three times in 24 hours. Only mechanical enforcement at every layer can deliver the discipline the project requires.

Propagation gate `CM-COVENANT-114-75-PROPAGATION` enforces this anchor literal across the ~44-file consumer fleet. Paired mutation strips the literal → gate FAILs.

Canonical authority: constitution submodule `Constitution.md` §11.4.75.

Non-compliance is a release blocker^{2026 05 20} regardless of context.

§ . . — Containers-submodule mandate
(User mandate, — —)

Forensic anchor — direct user mandate (verbatim, 2026-05-20):

"For any work or requirements of running services or codebase inside the Containers (Docker / Podman / Qemu / Emulators, and so on) we MUST USE / INCORPORATE the Containers Submodule properly: <https://github.com/vasic-digital/containers> (`git@github.com:vasic-digital/containers.git`). Containers Submodule contains all means for us to Containerize our code and services! If any feature or Containing System is missing or not supported we MUST EXTEND IT properly like we do all of our projects! All these details MUST BE added to our root constitution (constitution Submodule) - Constitution.md, AGENTS.md, CLAUDE.md, QWEN.md - and followed and respected! No bluff work is allowed of any kind!"

§-slot history note. This anchor was originally drafted as §11.4.75 but renumbered to §11.4.76 per the §11.4.71 fetch-before-push mandate after a concurrent upstream commit (`0a70083`) landed §11.4.75 (Mechanical Enforcement Without Exception). Same collision-resolution pattern as §11.4.75's own renumber-from-§11.4.74.

The mandate. For ANY containerized workload — Docker, Podman, Qemu, Kubernetes, container-backed emulators (Android emulator, VM-based testbeds, etc.) — every consuming project MUST:

1. **Use the canonical Containers Submodule.** `https://github.com/vasic-digital/containers` (`digital.vasic.containers`) is the authoritative library for runtime auto-detection, endpoint discovery, lifecycle/health management, compose orchestration, cross-build, emulator integration, and on-demand service boot. Mirrored across the four §2.1-canonical hosts (GitHub + GitLab + GitFlic + GitVerse).
2. **Install as a Git submodule.** Container-orchestration concerns enter the consuming project as a

```
git submodule add git@github.com:vasic-digital/containers.git  
<path>/containers
```

 . Project `go.mod` (or equivalent dependency manifest) consumes via `replace digital.vasic.containers => ./<path>/containers` during development; production builds resolve through pinned commit SHAs.

3. **Boot infrastructure on demand.** Tests, CLI doctor commands, and local-development workflows that depend on Postgres / Redis / OTel / message-brokers / emulators MUST invoke the Containers Submodule's `pkg/boot` + `pkg/compose` + `pkg/health` APIs to bring infra up automatically. Operators MUST NOT be required to manually start `podman machine` / `docker compose up` before tests — the boot is part of the test entry point. This is the **on-demand-infra invariant**.
4. **Extend the Containers Submodule, never reimplement.** Per §11.4.74's extend-don't-reimplement rule applied specifically to containerization: if a runtime (e.g., a new emulator type) or lifecycle primitive (e.g., a new health-check kind) is missing, the consuming project MUST add it to `vasic-digital/containers` (PR upstream + bump the consuming project's submodule pointer) — never as a parallel implementation inside the consuming project.
5. **Anti-bluff captured-evidence requirement.** Every integration test that claims to exercise a containerized component MUST actually boot that component via the Containers Submodule. Tests that pass without containers actually running (e.g., short-circuit fakes that bypass the boot path) are a §11.4 bluff violation. A passing test MUST imply the infra was up.
6. **Composition with §11.4.74.** The Containers Submodule itself is one of the Submodules surveyed by `Catalogue-Check`. A project tracker row touching containerization MUST record `Catalogue-Check: extend vasic-digital/containers@<sha>` (or `reuse`) — `no-match` is only valid if the consuming project demonstrates the Containers Submodule's API cannot model the workload (rare).

Anti-bluff captured-evidence gate (planned). `CM-CONTAINERS-USED` :

- For every PR that touches `Dockerfile*`, `*compose*.yaml`, `*podman*.yaml`, `Vagrantfile`, or any code under a directory matching `(test|integration|e2e)/.*infra` heuristic, the gate scans for an import of `digital.vasic.containers/...`. Absence in a non-trivial container-touching PR is a FAIL.
- The gate's paired mutation strips the import and asserts the gate FAILs.

Composition. Composes with §11.4.74 (catalogue-first discovery), §11.4.75 (mechanical enforcement — this clause IS one of the things mechanical enforcement protects), §3 (submodule propagation), §11.4.31 (Submodule-

Dependency-Manifest), §11.4.36 (`install_upstreams` on clone), §11.4.28 (Submodules-As-Equal-Codebase: the Containers Submodule itself follows all dev-cycle rules), §11.4.65 (multi-format export for docs maintained inside the Containers Submodule), §1.1 (paired-mutation gates protect the Containers Submodule's own boot / health primitives).

Why this matters. Without the on-demand-infra invariant, integration tests silently degrade to "pass without infra running" (a §11.4 bluff) or break with cryptic "connection refused" errors that operators waste hours debugging. The Containers Submodule turns infra-boot into a typed, programmable, anti-bluff-tested capability rather than a sequence of imperative shell commands. Consuming projects converge on a single shared lifecycle/health model instead of each project hand-rolling their own.

Classification: universal (per §11.4.17). Applies to every project that runs code inside any containerization or virtualization runtime.

Canonical authority: constitution submodule `Constitution.md` §11.4.76.

Non-compliance is a process violation; landing containerized infra without consuming the Containers Submodule (i.e., reinventing compose orchestration in-project) is a release blocker.

2026 05 20

§ . . — Regeneration-mechanism-
required mandate (User mandate,
— —)

Forensic anchor — direct user mandate (verbatim, 2026-05-20):

"We must be sure that after excluding anything from Git versioning we still have the mechanism which will out of the box obtain or re-generate missing content! Add this mandatory safety rule / constraint into ours root (constitution Submodule) `Constitution.md`, `CLAUDE.md`, `AGENTS.md`, `QWEN.md` or any other required document / file from the constitution Submodule! Do not forget to fetch and pull Submodule first! Commit and push all Submodule changes to all upstreams!"

Forensic incident (2026-05-20T15:00Z, parent project, a consuming project). The parent's audio Tier 1 `commit_all.sh` stalled four hours on `git add -A` while scanning 274 GiB of untracked `.git-backup-20260520T084610Z-pre-orphan-kill/` + 159 GiB `RKTools/linux/` + 167 GiB `qa-results/` content — all of which should never have been tracked. The natural remediation is to add those paths to `.gitignore`, BUT doing so without a regeneration mechanism would silently orphan fresh clones: a new operator running `git clone && cd <project>` would find missing RKTools, missing test infrastructure, missing build outputs they cannot reproduce. §11.4.77 codifies the universal rule that closes this gap.

The mandate. Every `.gitignore` entry that excludes either (a) more than ~100 MiB of content, OR (b) any artifact essential to building / running / testing the project, MUST be accompanied by a documented + automated mechanism that re-obtains or re-generates the excluded content on a fresh clone. The closed-set choices are:

1. **Re-obtain** — download from an authoritative source (vendor tarball, SDK installer, npm / pip / cargo / go-mod registry, container image registry, dedicated git submodule, S3, GCS, vendor-internal artifact server) via a script that runs deterministically on a fresh clone.
2. **Re-generate** — produce the content from tracked source code via a script (build pipeline, code generation, asset rendering, captured-evidence replay, container build, kernel build, image packaging) that runs deterministically on a fresh clone.

Required artefacts per qualifying `.gitignore` entry (each is independently mandatory):

1. **Metadata file at `.gitignore-meta/<entry-slug>.yaml`** (or equivalent project-canonical location declared in the project's `CLAUDE.md`) carrying: the gitignore-pattern verbatim, `mechanism-type: re-obtain | re-generate`, `script-path: <repo-relative path>` (the deterministic regeneration endpoint), `expected-disk-usage: <bytes-or-human-readable>`, `vendor-url-or-source: <URL or in-tree path>` when applicable, `integrity: { algorithm: sha256 | md5, value: <hex> }` when the content is fetchable from a known-stable mirror, `requires-network: true | false`, `requires-credentials: true | false` (and which credentials slot per §11.4.10).

2. **Entry in the post-clone bootstrap script** (`scripts/setup.sh` or the project-canonical equivalent). The script **MUST** run the regeneration mechanism non-interactively on a fresh clone or be invocable as an idempotent step (`scripts/setup.sh --regenerate <entry-slug>`).
3. **Pre-build gate** that verifies either the regenerated content is present at the expected path OR the regeneration script ran successfully (e.g., a stamp file under `.gitignore-meta/.regenerated/<entry-slug>.ok` with mtime $\geq$ source freshness window).
4. **Documentation update** — README plus the relevant `docs/guides/*.md` guide describing the mechanism, including the manual fallback (vendor URL, alternative download mirrors), the time + disk-usage budget, and per-§11.4.10 credentials requirements.

No escape hatch. Adding a bare `.gitignore` entry that excludes significant or essential content without the accompanying mechanism is itself a §11.4 PASS-bluff variant — the codebase appears complete to the casual eye, every pre-build / post-build gate goes green, but a fresh clone cannot build / run. There is no `--skip-regen-mechanism` , `--gitignore-is-enough` , `--operator-already-has-content` flag.

Composition with sister anchors:

- **§11.4.6 (no-guessing)** — the mechanism **MUST** be verified working end-to-end on a sandbox clone before the `.gitignore` entry lands. `LIKELY` works on a fresh clone is the exact bluff §11.4.6 forbids.
- **§11.4.65 (universal Markdown export)** — generated HTML / PDF siblings are an instance of `re-generate` ; the auto-regeneration helper (`sync_all_markdown_exports.sh`) is the §11.4.77 mechanism for that gitignored content where applicable.
- **§11.4.66 (interactive clarification)** — if the agent is uncertain whether a candidate `.gitignore` addition qualifies under §11.4.77, the agent **MUST** ASK the operator via the platform's interactive question mechanism rather than guess.
- **§11.4.71 (pre-push fetch + integrate)** — the regeneration mechanism's integrity hash and vendor URL **MUST** be re-validated against upstream before pushing the §11.4.77-governing commit; stale hashes are a §11.4.6 violation.

- **§11.4.74 (catalogue-first + extend-don't-reimplement)** — if the regeneration mechanism can be implemented by extending a Submodule already in `vasic-digital/HelixDevelopment` (e.g., a reusable downloader / decompressor helper), the project MUST extend rather than reimplement.
- **§11.4.75 (Mechanical Enforcement Without Exception)** — the local `pre-commit` hook (Layer 1) MUST refuse a commit that adds new `.gitignore` lines without a corresponding `.gitignore-meta/*.yaml`; the CI workflow (Layer 5) MUST replay the gate on every push. Paired meta-test mutation: strip a metadata YAML entry → gate FAILs.
- **§11.4.76 (Containers-submodule mandate)** — container images excluded from VCS are regenerated via `vasic-digital/containers pkg/boot` + Dockerfiles tracked in-tree; the `.gitignore-meta/<entry-slug>.yaml` references the Containers Submodule build path.
- **§9 / §9.2 (zero-risk data safety)** — the regeneration mechanism MUST be pre-tested in a sandbox clone with a hardlinked-backup safety net BEFORE relying on it for live operator clones; a regeneration script that silently corrupts content on a fresh clone is a §9 violation.
- **§3 (propagation order)** — submodules consuming this constitution submodule inherit §11.4.77; their own `.gitignore` entries are subject to the same mechanism requirement, propagated via their own `CLAUDE.md` / `AGENTS.md` / `QWEN.md`.

Anti-bluff captured-evidence gate (planned). `CM-GITIGNORE-REGEN-MECHANISM` :

- For every PR that adds a line to `.gitignore` (or to any `*.gitignore` file in a sub-tree), the gate scans for a matching entry in `.gitignore-meta/`, verifies the `script-path` exists and is executable, verifies the metadata YAML parses and carries the required keys, and verifies the post-clone bootstrap references the mechanism.
- Bare-line additions (e.g., `echo 'qa-results/' >> .gitignore`) without the metadata sibling are a FAIL.
- The gate's paired §1.1 mutation strips one required YAML key (e.g., `script-path:`) and asserts the gate FAILs.

Why this matters. Without §11.4.77, the natural `.gitignore` ergonomics — "this is too big to track, exclude it" — leads to fresh-clone orphanage: operators or CI runners (or new developers, or end-user mirror sites) clone the project and discover essential content is missing, with no documented way to obtain it. The cost is paid at every fresh-clone event, every CI bootstrap, every operator onboarding, every disaster-recovery exercise. §11.4.77 forces the cost to be paid once, at the `.gitignore` -addition commit, where it is bounded and visible.

Forensic anchor crosswalk to today's incident. Without §11.4.77, the natural fix to today's `commit_all.sh` 4-hour-stall — adding `.git-backup-*`, `RKTools/linux/`, `qa-results/` to `.gitignore` — would have orphaned every fresh clone of a consuming project. With §11.4.77, each of those entries now requires: (a) a `.gitignore-meta/<slug>.yaml` declaring the mechanism (e.g., `.git-backup-*` is a transient artefact safe to re-create per §9.2 hardlinked-backup; `RKTools/linux/` is re-obtainable from Rockchip's vendor tarball; `qa-results/` is re-generatable by re-running the on-device test cycle); (b) bootstrap-script entries; (c) gates; (d) docs. The bare `.gitignore` addition is now mechanically blocked.

Classification: universal (per §11.4.17). Applies to every project that excludes content from VCS for size / regenerability / vendor-provenance reasons.

Canonical authority: constitution submodule `Constitution.md` §11.4.77.

Non-compliance is a release blocker regardless of context.

2026 05 20

§ . . — CodeGraph code-intelligence
mandate (User mandate, — —)

Forensic anchor — direct user mandate (verbatim, 2026-05-20):

"All gathered knowledge and steps MUST BE added as future reusable experience into our constitution Submodule on setting up / configuring and using codegraph with all supported CLI agents we may use for our codebase! Make codegraph MANDATORY CHOICE for this purpose for all of our project and add proper comprehensive details about everything codegraph related into our root constitution (constitution Submodule) Constitution.md, AGENTS.md, CLAUDE.md, QWEN.md and other constitution Submodule relevant files! All project which do not have configured and installed codegraph yet MUST DO IT and MUST USE IT!"

The mandate. Every consuming project under this Constitution MUST install, initialize, and use **CodeGraph** (<https://github.com/colbymchenry/codegraph> , npm package `@colbymchenry/codegraph`) as its code-intelligence layer for AI coding agents. CodeGraph builds a local SQLite semantic knowledge graph of a codebase and exposes it to agents over the Model Context Protocol (MCP) — 100% local, no cloud, no external API (compatible with the Local-Only CI/CD posture). A project worked on by AI coding agents that has not yet installed and wired CodeGraph MUST do so; using it is not optional.

1. **Install.** `npm install -g @colbymchenry/codegraph` (Node.js 18+). The npm global prefix MUST be user-writable — installation MUST NOT require `sudo` / `su` (§9 / §12 host-safety + the no-privilege-escalation posture). Native modules (`better-sqlite3` , `tree-sitter`) compile on install with a WASM fallback.
2. **Initialize + index.** `codegraph init` at the project root creates `.codegraph/config.json` (tracked) and `.codegraph/codegraph.db` (a regenerable build artifact — gitignored per §11.4.30 / §11.4.77, with `codegraph index` as its declared regeneration mechanism). The `config.json` `exclude` list MUST exclude (a) other-owned submodules and vendored trees — index the consuming project's own domain code — and (b), non-negotiably, every credential/secret path (`.env*` , keystores, signing keys, service-account JSON) per §11.4.10. A secret reaching the index is a §11.4.10 violation. `codegraph index` builds the graph; `codegraph sync` keeps it fresh.

CRITICAL — filter mechanism. CodeGraph v1.x is zero-config: it uses BUILT-IN skip lists (`node_modules` , `vendor` , `dist` , `build` , `target` , `.env` , `Pods` , `.next`) PLUS the project root `.gitignore` to decide which files to index. The `.codegraph/config.json` `include` / `exclude` fields are **not used** for file filtering — all filtering happens through the ignore matcher seeded with the built-in defaults + root `.gitignore` patterns. To control what gets indexed, use root-anchored `.gitignore` patterns (start with `/` to match only the top-level directory) and `!` negations to re-include owned paths under a default-skipped parent (e.g., `!/vendor/widevine_*` re-includes owned vendor code that CodeGraph's built-in `vendor` skip would otherwise block).

MUST-EXCLUDE classes (actual rubbish — build outputs, caches, credentials, binaries):

```
** (a) Build outputs and regenerable artifacts: ** \out/, \prebuilts/
\, \rockdev/, \RKTools/, \compressed/.

** (b) Caches, dependencies, tooling: ** \.gradle/, \__pycache__/,
\.env/, \flash_logs/, \logs/.

** (c) Credentials, secrets, QA/recording artifacts: ** \secrets/,
\scripts/testing/secrets/, \**/*.env, \qa-results/, \recordings/.

** (d) Config.json include/exclude (DEPRECATED — not enforced by v1.x
engine): ** maintained as documentation-only for future compatibility.
```

ALL SOURCE CODE MUST REMAIN INDEXABLE, including AOSP platform trees (`frameworks/` , `external/` , `art/` , `cts/` , `bionic/` , `packages/` , `system/` , `kernel-5.10/` , `hardware/*/` , `vendor/` , `device/*/`) — these are tracked source and MUST NOT be gitignored for CodeGraph purposes. Only the rubbish classes above may be excluded. A project with 1M+ tracked files (AOSP fork) will produce a correspondingly large index (30-40 GB for the first init, with subsequent `codegraph sync` operations fast because they use `git diff`). The large initial size is expected and correct — the index represents the full codebase the AI agent needs to understand cross-references.

Universal mandatory-exclude baseline set (§11.4.78 + §11.4.79 refinement — binds EVERY consuming project). Independently of project, the CodeGraph index MUST NEVER include the following file-classes anywhere —

they are indexed-rubbish in every project; a consuming project MAY add its own project-specific rubbish paths on top per §11.4.35 but MUST NOT drop any baseline class: (a) **build outputs / regenerable artifacts** — `/out/` , `/build/` , `/dist/` , `/target/` , `bin/` , `obj/` , `.compressed/` , and any generator-produced tree; (b) **caches / dependency / tooling dirs** — `**/ccache/` , `**/.ccache/` , `.gradle/` , `__pycache__` , `.pytest_cache/` , `.mypy_cache/` , `.ruff_cache/` , `node_modules/` , `.next/` , `.cache/` , `.terraform/` , `.venv/` , `logs/` ; (c) **credentials / secrets** — `.env` , `.env.*` , `**/*.env` , `secrets/` , `**/*secret*` , `keystores` / `signing-keys` / `service-account JSON`, and every credential-bearing path per §11.4.10 (a secret reaching the index is a §11.4.10 violation, never negotiable); (d) **QA / recording / large regenerable corpora** — `**/qa-results/` , `**/recordings/` , and prebuilt-binary trees. Realized via root-anchored `.gitignore` patterns (CodeGraph v1.x zero-config filter = built-in skip list + root `.gitignore` , NOT `config.json` include/exclude). Own-org submodule sources stay INCLUDED per §11.4.79; third-party vendored trees EXCLUDED. This baseline is UNIVERSAL (§11.4.17) — dropping any class is a §11.4.78 / §11.4.79 violation, and no source-code tree may be excluded to satisfy it.

3. **Wire every supported CLI agent.** The CodeGraph MCP server (`codegraph serve --mcp` , stdio transport) MUST be registered with every AI coding CLI agent the project's developers use. Registration is project-scoped and committed where the agent supports it (e.g. Claude Code `.mcp.json` , OpenCode `opencode.json` , Qwen Code `.qwen/settings.json` , Crush `.crush.json`); host-local where the agent stores MCP config outside the repository (e.g. Kimi CLI `~/.kimi/mcp.json`). Every MCP config MUST reference the bare `codegraph` command resolved on `PATH` — never a hardcoded host path (the no-hardcoding posture). Claude Code is the canonical primary agent; the others MUST work too.
4. **Anti-bluff verification is mandatory.** CodeGraph integration MUST be covered by an anti-bluff verification suite (the canonical reference implementation is a consuming project's `scripts/verify-codegraph.sh` + `tests/codegraph/` , six layers): index reality, query correctness, MCP-protocol JSON-RPC, per-agent connectivity, per-agent end-to-end LLM drive, and a falsifiability rehearsal that mechanically proves the suite FAILs when the index is deliberately broken. The per-agent end-to-end layer MUST use an

unforgeable challenge — a fact obtainable only by calling a CodeGraph MCP tool (e.g. the index node count via `codegraph_status`) — so an agent answering from its own file-reading tools cannot produce a false PASS. An agent that genuinely cannot be driven end-to-end in the test environment (missing credentials, exhausted quota, environment incompatibility) is recorded as a documented SKIP gap per §11.4.3 — never a faked PASS.

5. **Comprehensive documentation.** Every project MUST carry a `docs/CODEGRAPH.md` (or canonical equivalent) describing install, initialization, indexing, per-agent wiring, the verification suite, and troubleshooting — kept in sync per §11.4.12 / §11.4.65.
6. **Catalogue-first.** CodeGraph is a third-party developer tool, not an owned submodule — it does NOT enter the project as a `git submodule` (per §11.4.74 it is consumed as the published npm package) and it does NOT add a Git remote.

Anti-bluff captured-evidence gate (planned). `CM-CODEGRAPH-WIRED` : for every consuming project, the gate verifies `.codegraph/config.json` exists with the §11.4.10 secret-exclusions present, every developer-used agent carries a CodeGraph MCP registration, and the verification suite exists and is executable. The paired §1.1 mutation removes a secret-exclusion from `config.json` and asserts the gate FAILs.

Composition. Composes with §11.4.3 (per-environment-topology SKIP for un-runnable agents), §11.4.10 (credentials never indexed), §11.4.12 + §11.4.65 (`CODEGRAPH.md` kept in sync + exported), §11.4.30 / §11.4.77 (the `.codegraph/codegraph.db` artifact is gitignored with `codegraph_index` as its regeneration mechanism), §11.4.74 (catalogue-first — CodeGraph is consumed, not reimplemented), §11.4 (the verification suite is anti-bluff by construction — CI green is necessary, never sufficient), §1.1 (paired-mutation gate).

Why this matters. AI coding agents that explore a codebase by repeated file scanning consume tokens and tool calls wastefully and build a shallow, drift-prone mental model. A pre-indexed semantic graph gives every agent instant, consistent symbol / caller / callee / impact resolution. Standardising on one tool — CodeGraph — across every Helix project means the capability is wired once, verified anti-bluff once, and reused everywhere, instead of each project hand-rolling ad-hoc context tooling.

Classification: universal (per §11.4.17). Applies to every project that is worked on by AI coding CLI agents.

Canonical authority: constitution submodule `Constitution.md` §11.4.78.

Non-compliance is a process violation; a project worked on by AI agents without CodeGraph installed, wired, and anti-bluff-verified is in breach of this mandate.

§ 11.4.79
2026-05-21
– Own-org submodules MUST be included in the CodeGraph index (User mandate, –)

Forensic anchor — direct user mandate (verbatim, 2026-05-21):

"All Submodules we use in the project and that are part of organizations to which we have the full access via GitHub, GitLab and other CLIs MUST BE included into the codegraph database and initialized / scanned / synced!"

The mandate. This extends §11.4.78 step 2's exclude-list refinement with a per-submodule-ownership split. Every consuming project's `.codegraph/config.json` MUST distinguish:

Submodule class	Treatment in CodeGraph index
Own-org — full write access via the project's CLIs (canonical orgs: <code>vasic-digital</code> on GitHub/GitLab/GitFlic/GitVerse + <code>HelixDevelopment</code> on GitHub)	MUST be INCLUDED. Per-submodule sources MUST be indexed so AI coding agents resolve symbols, callers, and impact across the whole repo graph — not just the consuming project's domain code.
Third-party — no write access, vendored only (the §11.4.74 <code>no-match</code> → <code>vendor</code> path, e.g. an external SDK like <code>gopkg.in/telebot.v3</code>)	MUST be EXCLUDED. Indexing third-party code wastes the local index budget and creates symbol-noise the project's contributors cannot fix.

This refines, NOT contradicts, §11.4.78 step 2 (which spoke of "other-owned submodules" generically). The classifier is **write-access via the project's own CLIs**, not "internal vs external" subjectively — own-org submodules are the project's own code under another path; third-party submodules are upstream code Herald extends without owning.

Operational steps for every consuming project:

1. **Fetch + pull latest of every submodule** before re-indexing — `git submodule update --remote --merge` from the project root, then verify the project still builds. Third-party submodule pins (e.g. `gopkg.in/telebot.v3` pinned at `v3.3.8` to keep the Go import path stable) MUST be respected and rolled back if a `--remote` advances past the load-bearing pin.
2. **Adjust `.codegraph/config.json` `exclude` to keep own-org submodules in scope.** Third-party submodule paths MUST be explicitly listed in `exclude` (per §11.4.78 step 2's per-credential exclusion principle); own-org submodule paths MUST NOT appear in `exclude`.
3. **Re-index.** Run the project's canonical CodeGraph initialiser (e.g. `scripts/codegraph_setup.sh`) so `codegraph index` (or `codegraph sync`) rebuilds the index with the corrected exclude list.
4. **Verify symbols across own-org submodules.** The project's `scripts/codegraph_validate.sh` (per §11.4.78 anti-bluff verification) MUST probe at least one symbol that lives ONLY inside an own-org submodule (e.g. `bg.TaskQueue` from `submodules/background/interfaces.go`). A successful query proves the include path actually reached the submodule's sources — not a §107 PASS-bluff where validate runs against a stale index.
5. **Paired §1.1 mutation:** temporarily add the own-org submodule path to `exclude`, re-index, run validate — it MUST FAIL on the cross-submodule probe. Restore. This mutation proves the validate is not bluffing about the include reach.

Composition. Composes with:

- §11.4.74 (catalogue-first): the `own-org` classifier matches the `reuse/extend` catalogue-check disposition; `third-party` matches `no-match` → `vendor`.
- §11.4.78 (CodeGraph mandate): this §11.4.79 refines step 2's exclude-list contract without weakening any other step.

- §11.4.10 (credentials never tracked): the exclude list MUST still cover every `.env*`, keystore, signing key, service-account JSON regardless of submodule ownership.
- §1.1 (paired mutation): every validate probe added per step 5 above MUST land with a mutation pair.
- §107 (end-user usability): an indexed graph that lies about which symbols are reachable is a §107 PASS-bluff against the AI coding agents that consume it.

Why this matters. When AI coding agents work across a project that vendors own-org capability modules (e.g. `commons_constitution` consuming the Helix-stack 9-module surface under `submodules/`), excluding those modules from the index forces the agents back into shallow file-scan exploration of half the codebase — exactly the failure mode §11.4.78 exists to prevent. The agents lose call-graph resolution across the seam between domain code and own-org infra code. Including own-org submodules is the difference between agents seeing the project as a unified graph and agents tripping over an invisible boundary at the submodule directory.

Classification: universal (per §11.4.17). Applies to every project under this Constitution that has CodeGraph configured and own-org submodules vendored.

Canonical authority: constitution submodule `Constitution.md` §11.4.79.

Non-compliance is a process violation; an AI-agent-worked project whose own-org submodules are excluded from CodeGraph (when they could be included) is in breach of this mandate. Severe cases (own-org submodules silently excluded ¹¹⁴⁸⁰ WITHOUT an audit trail in `.codegraph/config.json` comments) are release blockers.

2026 05 21

§ . . — CodeGraph regular-update +
sync automation mandate (User mandate,
— —)

Forensic anchor — direct user mandate (verbatim, 2026-05-21):

"We MUST regularly check for the updates and execute codegraph npm updates so the latest version of it is always installed on the host machine! After update executes with success use the following commands (and validate and verify them using codegraph help) to sync all: [list of codegraph commands]. Make sure we have proper full automation bash scripts which will run regularly and that these are part of the constitution Submodule since they MUST BE available to all projects which do respect and follow our constitution rules and mandatory constraints! Make sure all updates, sync processes we do and important codegraph related events are all documented under docs/codegraph in Status and Status_Summary documents (another area / context to regularly update and sync) and regularly export them like all other Status docs into the PDF and HTML!"

The mandate. Every consuming project under this Constitution MUST regularly check for CodeGraph npm-package updates, install the latest stable version, and re-sync its local index. The automation lives in the constitution submodule itself so every consuming project gets it via the inherited tree without reimplementation.

Three deliverables:

1. **<constitution>/scripts/codegraph_update.sh** — checks the installed `@colbymchenry/codegraph` version against npm's latest, runs `npm update -g @colbymchenry/codegraph` (or `npm install -g @colbymchenry/codegraph@latest`) if a newer version exists, captures the old + new version + timestamp into the constitution's `docs/codegraph/Status.md` audit trail, then exits 0 only when `codegraph --version` produces the expected post-update output. Idempotent: re-runs on the same day are no-ops. Anti-bluff (§107): the "update succeeded" PASS MUST observe the new version, not just `npm update` exit code.
2. **<constitution>/scripts/codegraph_sync.sh** — after a successful update, runs the canonical sync sequence inside the consuming project (passed as an argument or auto-resolved via `find_constitution.sh` parent walk):
`codegraph status` (baseline) → `codegraph sync .` (incremental update) → `codegraph status` (post-sync) → `scripts/codegraph_validate.sh` (the project's anti-bluff verifier, per §11.4.78 step 4). Each step's output is appended to the project's `docs/codegraph/Status.md` ledger. The sync

MUST consult `codegraph help` BEFORE invoking any command to verify the subcommand still exists with the same shape (tools evolve). The canonical sync sequence covered by `codegraph help` :

```
codegraph                # Run interactive installer
codegraph install        # Run installer (explicit)
codegraph init [path]    # Initialize in a project (--index to
also index)
codegraph uninit [path]  # Remove CodeGraph from a project (--
force to skip prompt)
codegraph index [path]   # Full index (--force to re-index, --
quiet for less output)
codegraph sync [path]    # Incremental update
codegraph status [path]  # Show statistics
codegraph query <search> # Search symbols (--kind, --limit, --
json)
codegraph files [path]   # Show file structure (--format, --
filter, --max-depth, --json)
codegraph context <task> # Build context for AI (--format, --
max-nodes)
codegraph affected [files...] # Find test files affected by changes
codegraph serve --mcp     # Start MCP server
```

3. `<constitution>/docs/codegraph/Status.md + docs/codegraph/Status_Summary.md` — append-only ledgers tracking every CodeGraph-related event across the consuming project: install events, npm-update events, sync runs, index regenerations, validate runs (with PASS/FAIL counts), HRDs opened/closed against CodeGraph. Both ledgers MUST be exported to `.html` and `.pdf` siblings on every edit (per §11.4.65 multi-format export mandate). The `Status_Summary.md` is the operator-readable derived index (per §11.4.53 fixed-summary backfill convention).

Operational cadence. The automation MUST run at least weekly (per §11.4.45 status-digest cadence). Consuming projects MAY wire it more frequently (e.g. daily via cron or per-CI-run) but never less. Cadence is not contractually fixed because environments differ (development workstation vs CI vs production-locked host); the floor is "weekly" because slower than that risks accumulated drift.

Scripts are inherited by reference. Per §3 submodule inheritance, consuming projects do NOT copy `codegraph_update.sh` / `codegraph_sync.sh` into their own `scripts/` directory. They invoke the constitution submodule's scripts directly

(e.g. `bash ${CONST_DIR}/scripts/codegraph_update.sh && bash ${CONST_DIR}/scripts/codegraph_sync.sh .`). This is the single-source-of-truth pattern — when the constitution updates the scripts, every consuming project picks up the change on its next submodule update.

Composition. Composes with §11.4.78 (CodeGraph parent mandate), §11.4.79 (own-org submodule inclusion), §11.4.10 (credential exclusions stay applied), §11.4.45 (status-digest cadence), §11.4.53 (Fixed_Summary backfill), §11.4.65 (multi-format export), §107 (each sync run's "ok" MUST carry observed-version evidence, not just exit code), §1.1 (paired mutation: temporarily downgrade installed codegraph version → script MUST detect drift and self-correct → restore).

Why this matters. A CodeGraph index built against an older version of the indexer may miss symbols that newer versions resolve correctly. Stale indexes silently lie to AI agents about what symbols exist and how they connect. Regular automated updates + sync close the loop: tooling stays current, the index reflects current source, AI agents see the truth. Manual "run codegraph index when I remember" is exactly the failure mode this mandate forbids.

Classification: universal (per §11.4.17). Applies to every project under this Constitution that has CodeGraph configured per §11.4.78.

Canonical authority: constitution submodule `Constitution.md` §11.4.80.

Non-compliance is a process violation; severe cases (consuming project has not run `codegraph_update.sh` in >2 weeks AND has open AI-agent work) are release blockers — the agents' index reflects a stale tool version and may produce silently wrong reasoning.

§ . . — Cross-platform-parity mandate
(User mandate, — —)

Forensic anchor — direct user mandate (verbatim, 2026-05-21):

"Any Linux-only blocker / issue we have MUST BE created macOS and other supported platforms equivalent! So, depending on platform proper implementation will be used for particular OS! EVERYTHING MUST BE PROPERLY EXTENDED AND UPDATED!"

The mandate. Every consuming project whose `supported-platforms` manifest lists more than one OS MUST, for every feature/test/gate/challenge/mutation that depends on platform-specific primitives, ship a per-OS-equivalent implementation chosen at runtime via `uname -s` (or equivalent platform detection — `[[ $OS == 'Windows_NT' ]]`, `sw_vers -productName`, `/etc/os-release` on Linux distros, etc). A Linux-only gate that ships without macOS / Windows / BSD equivalents (when the kernel of those OSes permits one) is a §11.4 / §107 PASS-bluff at the platform-coverage layer: the test reports GREEN on Linux but tells nothing about whether the feature works for end users on every supported platform.

Three sub-mandates.

(A) Per-OS implementation REQUIRED. Every feature whose Linux implementation uses a `cgroup` / `systemd` / `/proc` primitive MUST have a documented per-OS equivalent (POSIX `setrlimit` / `ulimit`, macOS `launchd` `plist HardResourceLimits`, BSD `rctl`, Windows Job Object, etc) chosen via runtime platform detection. The wrapper/library code MUST dispatch to the correct branch with NO operator awareness required — the same operator-path invocation works on every supported platform. Canonical example from

```
vasic-digital/tmux : tmx new -s NAME dispatches to systemd-run --user  
--scope on Linux and to scripts/tmx-rlimit-wrapper.sh (POSIX  
setrlimit) on Darwin — same operator command, OS-correct mechanism per  
session.
```

(B) Per-OS tests REQUIRED. Every gate test that exercises platform-dependent behavior MUST have a per-OS branch under `case "$(uname -s)" in`. Each branch MUST exercise the platform's actual native primitive with positive captured evidence per §11.4.2 + §11.4.5 (e.g. read `ulimit -t` inside the session via `send-keys + capture-pane` on Darwin to verify `RLIMIT_CPU` applied; read `/sys/fs/cgroup/.../memory.max` content on Linux to verify `cgroup MemoryMax` applied). Platform-conditional `SKIP-with-reason` is acceptable ONLY when the platform genuinely cannot enforce the invariant at all (kernel limit — see (C) below), AND the SKIP message MUST cite the precise kernel/system limitation with a reproducer + a link to the project's honest-gap section in its operator guide (typically `docs/guide/README.md` or equivalent).

(C) Honest kernel-gap citation + adjacent equivalent test REQUIRED. Where a Linux primitive has NO macOS / other-OS equivalent due to a documented kernel/ OS limitation (canonical example: XNU does NOT enforce `RLIMIT_AS` /

`RLIMIT_DATA` / `RLIMIT_RSS` for unprivileged processes — verified by reproducer in `vasic-digital/tmux docs/guide/README.md §5.6`), the test MUST: (1) detect the gap at runtime, (2) `SKIP` with the exact kernel reason + reproducer + link to the project's honest-gap doc, (3) **provide an ADJACENT test that exercises the closest invariant the platform CAN enforce** (e.g. `RLIMIT_CPU` + `SIGXCPU` as the macOS proxy for "process is bounded under load"). The adjacent test MUST itself be anti-bluff per §11.4 with positive captured evidence and a paired §1.1 mutation.

Composes with. §11.4.1 (FAIL-bluffs equally forbidden — a Linux-PASS / Darwin-SKIP-without-honest-reason is a §11.4.81 violation), §11.4.2 (recorded-evidence — each branch's captured evidence is platform-specific), §11.4.3 (per-host-topology dispatch — §11.4.81 strictens §11.4.3 by requiring an equivalent test when the platform CAN enforce, not just SKIP-with-reason for genuine kernel gaps), §11.4.5 (captured evidence quality — platform-specific quality analysis), §11.4.4 (test-interrupt-on-discovery — discovering a Linux-only test without a Darwin equivalent triggers §11.4.81 mid-cycle), §11.4.6 (no-guessing — "platform X probably can't do Y" is a §11.4.6 violation; prove via reproducer or mark `UNCONFIRMED`), §11.4.20 / §11.4.70 (subagent-driven — cross-platform branches are independent enough to dispatch as parallel subagent work per branch), §11.4.27 (no-fakes-beyond-unit + 100% test-type coverage — platform branches inherit the 100% type-coverage discipline), §11.4.69 (universal sink-side positive-evidence taxonomy — every platform branch must produce a taxonomy-mapped sink-side evidence artefact), §107 (end-user usability — multi-platform projects whose tests cover only one platform deliver no end-user guarantee on the others).

Per-OS implementation equivalence catalogue (canonical examples — not exhaustive).

Linux primitive	Darwin equivalent	BSD equivalent	Windows equivalent
<code>systemd-run --user --scope (transient cgroup)</code>	<code>scripts/<rlimit-wrapper>.sh</code> using <code>ulimit -t (RLIMIT_CPU) + ulimit -u (RLIMIT_NPROC)</code>	<code>rctl process:0:deny:vmemory-use (root); ulimit -v (unprivileged limited)</code>	Windows Job Object via <code>CreateJobObjectW + SetInformationJobObject</code>
<code>cgroup MemoryMax</code> enforcement	XNU GAP — unprivileged <code>RLIMIT_AS</code> returns <code>EINVAL</code> ; root-only via <code>launchd HardResourceLimits</code> plist. Adjacent test: <code>RLIMIT_CPU + SIGXCPU</code> proxy	<code>rctl process:0:deny:memory-use (root)</code>	Job Object <code>JOB_OBJECT_LIMIT_PROCESS_MEMORY</code>
<code>cgroup TasksMax</code>	<code>RLIMIT_NPROC</code> via <code>ulimit -u (per-user, kernel-enforced)</code>	<code>rctl user:NN:deny:maxproc</code>	Job Object <code>JOB_OBJECT_LIMIT_ACTIVE_PROCESS</code>

Linux primitive	Darwin equivalent	BSD equivalent	Windows equivalent
<code>/proc/<pid>/oom_score_adj</code>	No equivalent — Darwin/BSD have no kernel OOM killer (different memory model). SKIP-with-reason; adjacent: rlimit bounds rogue session reach	<code>/dev/oom</code> -style not present in stock BSD	None at OS level (process governance via Job Object)
<code>/sys/fs/cgroup/... introspection</code>	<code>ulimit -a</code> readback inside session + <code>ps</code> ancestry	<code>procstat -r <pid></code>	<code>QueryInformationJobObject</code>
<code>kill -KILL</code> on cgroup process to test scope isolation	<code>kill -KILL</code> on tmux server pid directly; verify sibling servers survive via direct socket query	same	<code>TerminateProcess</code> + per-job survival check

Operational discipline. When a consuming project adds a new feature whose Linux path lands first, the SAME pull request MUST add the per-OS branches for every other platform in the project's `supported-platforms` manifest (or explicitly classify the platform per (C) — honest gap + adjacent test). Skipping the per-OS branches "for later" is itself a §11.4.81 violation: the gap accrues silently and the project ships Linux-PASS / other-OS-untested code while claiming multi-platform support.

Pre-build gate `CM-CROSS-PLATFORM-PARITY` (when implemented in the consuming project): scans every gate test for `case "$(uname -s)"` blocks, asserts a non-SKIP branch (or honest-gap citation per (C)) exists for each platform in the supported-platforms manifest. Paired §1.1 mutation: strip the Darwin branch from any one test → gate FAILs.

Classification: universal (per §11.4.17). Applies to every project under this Constitution whose supported-platforms manifest lists more than one OS. Single-platform projects are unaffected (no parity gap exists).

Canonical authority: constitution submodule `Constitution.md` §11.4.81.

Non-compliance is a release blocker on multi-platform projects. No escape hatch — no `--linux-only-acceptable`, `--no-platform-parity-required`, `--platform-gap-just-skip` flag exists. A Linux-PASS / other-OS-untested feature in a multi-platform project is a §11.4 / §107 PASS-bluff at the platform-coverage layer.

2026 05 22

§ . . — Iteration-speedup discipline mandate (User mandate, — —)

Forensic anchor — direct user mandate (verbatim, 2026-05-22):

"How can we speed-up this whole development and fixing process? ... Do not forget to all speed optimizations critical rules and mandatory constraints MUST BE all added into our root (constitution Submodule) `Constitution.md`, `CLAUDE.md`, `AGENTS.md` and `QWEN.md` and all other relevant constitution Submodules files!"

Iteration cycle time is a first-order quality enabler. A 40-min rebuild cycle that catches one defect per cycle bounds the project's defect-discovery rate at 1 defect per 40 min; a 15-min cycle bounds it at ~3× that rate. Slow cycles are a §11.4 PASS-bluff at the velocity layer — the project ships fewer validated features per unit of operator time than it should.

The 2026-05-22 forensic session that produced this anchor witnessed (a) 2 subagent watchdog stalls each ~10 min wasted, (b) 1 stalled `commit_all.sh` ~10 min lost on submodule cascade prompt, (c) 1 wrong-call-site speculative kernel

patch (§GQ.2) burning a full 45 min rebuild cycle before forensic evidence pointed to the correct site (§GT Fix B-1), (d) 1 stale `git index.lock` from a killed process blocking the next commit, (e) 1 auto- `git gc` maintenance colliding with urgent commit. Each line item cost between 5 and 45 minutes of wall-clock that should have been productive iteration time. The aggregate impact for that single morning was ~3 hours of operator wait dilated onto a job whose intrinsic compute was ~90 minutes.

The mandate. Every consuming project's build / test / commit / debug pipeline MUST adopt the following speedup disciplines AS MANDATORY. Each is independently enforceable:

(A) Phase forensic before any speculative source patch

Before applying ANY non-trivial source patch (kernel, framework, native code, HAL, ASoC), the agent MUST first complete Phase 1 of `superpowers:systematic-debugging` — read the error/log/call-chain end-to-end and identify a FACT-grade root cause. Speculative patches without Phase 1 evidence are §11.4.6 (no-guessing) violations AND §11.4.82 violations. A wrong-call-site rebuild cycle wastes 30-90 min of compute; the alternative (15-30 min of focused source reading) is a net positive every time. Phase 1 is mandatory regardless of operator time pressure — "we don't have time for forensics, just rebuild" is the §11.4.82 anti-pattern.

(B) Live-ADB-First (or live-equivalent) before any rebuild

Per §11.4.51 — strengthened by §11.4.82 to a release-blocker mandate. ~70% of audio / UI / boot-script / config changes are LIVE_ADB_TESTABLE (push + `setprop` + `pm install -r` + 30-second validation cycle). Only kernel / framework Java/AIDL / native C++-in-APEX / sepolicy / init.rc / ro.* / Android.bp / XML-resource-overlays / codec-XML-in-APEX changes are genuinely REQUIRES_REBUILD. Skipping the live-probe step for a LIVE_ADB_TESTABLE change is a §11.4.82 violation — costs operator 45 min of rebuild for what could have been 30 seconds of `adb push` .

(C) Pre-flight before launching rebuild orchestrators

Before any `nohup bash scripts/<rebuild>.sh &` invocation, the agent MUST run a ~30-second pre-flight that verifies (i) target device(s) reachable via ADB, (ii) downstream sinks (Arvus AVR, HDMI sink, Bluetooth peers) reachable as the test expects, (iii) host has sufficient memory and disk per §12.6 budget, (iv) no stale lock files in `.git/` or build directories, (v) no orphan git/build processes from prior killed iterations. A rebuild orchestrator launched against a broken precondition wastes 45 min discovering at the test stage what 30 seconds of pre-flight would have caught.

(D) Persistent build caches outside containers

Containerized AOSP / kernel builds MUST persist `ccache` AND any equivalent build-system intermediate cache (Soong cache, sccache, Gradle daemon state) on the host via bind-mount, NOT inside the ephemeral container layer. Without this, every container restart drops the cache and forces a cold rebuild (~40 min for AOSP+kernel on a fast workstation). With it, incremental rebuilds where only one driver / one app changed drop to 5-15 min (`ccache` hit rates 80-95%). This is a one-time setup of ~15 min for ~25-32 min savings per subsequent cycle.

(E) Module-only rebuild for loadable-module-only changes

When a patch touches ONLY one or more files inside `kernel/.../drivers/*/foo.c` AND the affected drivers are built as loadable modules (`CONFIG_*=m` in the active defconfig), the rebuild MUST be `make -C kernel M=<driver-path> modules` (~2-5 min) NOT a full kernel rebuild (~25 min). A

`flash_kernel_only.sh` or equivalent helper writes the rebuilt module(s) to the running device without re-flashing `system.img`. This saves 15-20 min per cycle for the ~10-30% of kernel patches that qualify. Built-in drivers (`CONFIG_*=y`) do not qualify — full rebuild required for those.

(F) Parallel multi-device testing

When the project owns more than one validation device (D3, D4, ... in topology terms), the test orchestrator MUST run autonomous validation cycles on EVERY device in parallel via background processes, with separate `qa-results/<TS>/<device-tag>/` output directories per device. This catches per-device topology

defects (per §11.4.3) one cycle earlier instead of waiting for a device-N-focused re-cycle. Direct wall-clock savings = 0 (parallel); amortized savings = 5-15 min per cycle when a device-specific defect surfaces.

(G) Subagent scope discipline + worktree isolation

Subagents dispatched per `superpowers:subagent-driven-development` MUST be:

- **Scope-bounded** to ≤ 30 min of focused work with intermediate output emitted as the subagent progresses (the canonical watchdog stalls at 600s of no-output; if a subagent writes its findings file incrementally it survives even when blocked on an external dependency).
- **Worktree-isolated by default** via the Agent tool's `isolation: "worktree"` parameter when supported, so parallel subagents do not contend on the same working tree and so a subagent crash never corrupts the conductor's state.
- **Single-responsibility** — one investigation OR one fix-implementation OR one doc draft, not multiple. The empirical pattern from the §GT investigation is that 8-task subagents stall while 2-3-task subagents complete reliably.

(H) Lock-file + stale-process hygiene

Killed or crashed agents leave behind `.git/index.lock`, `.git/.commit_all.lock`, `.lock` files in build directories, and orphan child processes (`git add -A`, `git pack-objects`, `git gc`). The next agent MUST detect these on session start (or after operator-visible crash) and clean them — never wait silently for a never-arriving release. Auto `git-gc` MUST be disabled in repos with concurrent multi-agent work (`git config gc.auto 0`) — its background invocation is a guaranteed lock-contention source.

(I) Cycle telemetry per §11.4.24 build-resource stats

Every iteration cycle MUST record (i) commit hash of the source under test, (ii) wall-clock of build + flash + test phases separately, (iii) the speedup-discipline flag set (which of A-H applied), (iv) outcome (PASS / FAIL / partial / stall). Aggregated weekly, this exposes which speedup gives the biggest empirical return on the

specific project, and which is being neglected. The §11.4.24 build-resource stats tracker is the canonical telemetry pipeline; §11.4.82 extends its row schema with iteration-discipline columns.

Composes with. §11.4.4 (test-interrupt-on-discovery + retest), §11.4.6 (no-guessing — speculative patches forbidden), §11.4.9 (batch-source-fixes-before-rebuild), §11.4.20 / §11.4.70 (subagent-driven), §11.4.24 (build-resource stats), §11.4.42 (iteration-discipline conductor loop), §11.4.43 (TDD-fix discipline — RED-test enables faster cycle reads), §11.4.50 (deterministic consistency — caches don't break determinism if ccache config disables time-sensitive headers), §11.4.51 (LIVE-ADB-first maximisation — strengthened to mandate here), §11.4.52 (autonomous-validation — every cycle's verdict is automation-produced), §11.4.58 (parallel-development PWU — parallel devices and parallel subagents are §11.4.58 instances), §12.7 (-j2 host-safety cap remains — speedups MUST stay within memory budget), §107 (end-user usability — fast iteration = fewer cycles → fewer defects shipped).

Pre-build gate. CM-ITERATION-SPEEDUP-DISCIPLINE (when implemented in consuming project): audits the most recent N iteration cycles for telemetry entries citing which of (A)-(I) applied; flags any cycle where none of (A)-(I) was active as a candidate for review. Paired §1.1 meta-test mutation: strip the speedup-flag column from a synthetic telemetry row → gate FAILs.

Operational discipline. When a consuming project ships an iteration speedup (e.g. lands a persistent-cache bind-mount, a parallel-device orchestrator, a module-only rebuild helper), the change MUST be classified as universal (per §11.4.17 — most are) OR project-specific (rare — e.g. a script that hardcodes D3 / D4 / Arvus IP 192.168.4.185). Universal speedups belong in this constitution submodule's helper inventory (or in a separate Containers submodule per §11.4.76); project-specific wrappers stay in the consuming repo.

Classification: universal (per §11.4.17). Applies to every project under this Constitution. The disciplines are language-/platform-/build-system-agnostic.

Canonical authority: constitution submodule Constitution.md §11.4.82.

Non-compliance is a release blocker. No escape hatch — no `--skip-phase1-forensic`, `--no-pre-flight`, `--rebuild-everything-always`, `--unlimited-subagent-scope`, `--ignore-locks`, `--no-telemetry` flag exists. The disciplines exist because their absence has been forensically demonstrated to cost the operator hours per session.

2026 05 22

§ . . — docs/qa/ end-user evidence mandate (User mandate, — —)

The mandate. Every feature that ships MUST carry a recorded end-to-end communication transcript plus any attached materials (screenshots, request/response payloads, audio, file uploads) committed under `docs/qa/<run-id>/` — one directory per feature run. A feature with no QA transcript is itself a §11.4 / §107 PASS-bluff: it claims to work but carries no auditable runtime evidence that an end user actually exercised it through the same interface they will use in production.

Forensic anchor (verbatim user mandate, 2026-05-22):

"every feature that ships MUST carry a recorded e2e communication transcript + any attached materials under `docs/qa/<run-id>/` (per-feature subdirectories). A feature with no QA transcript is itself a §107 PASS-bluff — it claims to work but has no auditable runtime evidence. Bot-driven automation MUST preserve full bidirectional communication threads as proof."

Operative rule.

1. Every consuming project MUST maintain a `docs/qa/` tree (no exception, no escape hatch). Each new feature run lands under `docs/qa/<run-id>/` where `<run-id>` is monotonic + greppable (timestamp, HRD-NNN, ATM-NNN, or other workable-item identifier — see §11.4.54).
2. The transcript MUST be full bidirectional — every prompt/command sent + every response received + every error message + every state change observed. One-sided ("we sent X") is not a transcript; both halves are required.

3. Attached materials MUST be committed alongside (screenshots in `.png` , payloads in `.json` / `.toon` / `.txt` , audio in `.wav` / `.mp3` , etc.). External-only links (Slack URL, Drive URL) are §11.4.13 sink-side violations — the evidence MUST live in-repo.
4. Bot-driven / agent-driven automation (a consuming project's QA bot, e.g. Herald's planned `qaherald` , or analogous binaries in other projects) MUST preserve the full conversation thread as the proof artefact. A bot that runs the round-trip but stores only the final PASS/FAIL line is itself a §107 bluff at the QA-automation layer.
5. CI / release gates MUST refuse to tag a version that has any feature-shipping commit without its matching `docs/qa/<run-id>/` directory present and non-empty. The `release` script in the consuming project enforces this; the constitution submodule's `scripts/qa_coverage_audit.sh` (when implemented) gives the universal scanner.

Composes with §11.4.2 (recorded-evidence requirement — `docs/qa/` is the canonical capture location for end-user-facing features), §11.4.5 (captured-evidence quality analysis applies per-transcript), §11.4.13 (sink-side evidence stays in-repo, not behind external links), §11.4.65 (multi-format export — transcripts may be exported to PDF/HTML siblings when operator-readable), §11.4.69 (universal sink-side positive-evidence taxonomy — `docs/qa/` is the "user-facing-channel" branch), §107 (an end-user feature without an end-user-channel transcript is a §107 PASS-bluff), §1.1 (paired mutation: delete a `docs/qa/<run-id>/` → release gate FAILs).

Classification: universal (per §11.4.17). Every project under this Constitution carries this rule. Single-feature pre-implementation projects are exempt until their first user-visible feature lands.

Canonical authority: constitution submodule `Constitution.md` §11.4.83.

Non-compliance is a release blocker. No `--qa-evidence-optional` , `--qa-transcript-later` , `--qa-bot-summary-suffices` flag exists.

§ 1.1 – Working-tree quiescence rule for subagent commits (User mandate, 2026-05-22)

Short tag: `working-tree quiescence` .

The mandate. No subagent commit may proceed while any concurrent mutation gate, paired-mutation experiment, or other in-flight mutation is live in the same checkout. Before `git add` , the committing agent MUST `grep` its own working tree for mutation markers (`MUTATED` for paired , `// always pass` , `return json.Marshal` shortcut paths, `// MUTATION / # MUTATION` annotations, `_mutated_*` filename suffixes, etc.) and explicitly account for every modified file in the staging area. Any unexplained file in the staging area triggers ABORT.

Forensic anchor (verbatim user mandate, 2026-05-22):

"no subagent commit may proceed while any concurrent mutation gate is in flight in the same checkout. Before `git add` , the committing agent MUST `grep` its own working tree for mutation markers (`MUTATED` for paired , `// always pass` , `return json.Marshal` shortcut paths, etc.). Any unexplained file in the staging area triggers ABORT."

Lesson (forensic case study). A consuming project's logo-fix subagent (Herald commit `72e81ab` , 2026-05-21) ran in a checkout where a paired §1.1 mutation gate had temporarily introduced an `// always pass` shortcut into `commons_auth/middleware.go` (JWT-bypass mutation, intended to be reverted by the same gate). The subagent's `git add` + `git commit` swept the mutation residue into the same commit as the unrelated logo fix, and the resulting commit was pushed to all four mirrors before any other agent caught it. The fix (Herald `d5bd360` , "SECURITY FIX: restore commons_auth/middleware.go JWT verify") landed within the hour, but the window during which production-equivalent binaries shipped with a bypassed JWT verify is a real security-defect window — small but non-zero — and is the canonical example of why this rule exists.

Operative rule.

1. Subagent (and main-thread) commit flows MUST run a pre- `git add` quiescence check that:
 - Greps the working tree for mutation markers (canonical list above; the consuming project MAY extend with project-specific markers).
 - Lists every modified / untracked / staged file and matches each one against the subagent's declared scope. Any file outside the declared scope → ABORT.
 - Cross-checks `git status --porcelain` against the subagent's task description; unaccounted entries → ABORT.
2. Any active mutation gate (paired-mutation experiment in progress, AVR/AVI mutation cycle, manual `// always pass patch`, etc.) MUST be serialised — the gate runs to completion (mutate → assert FAIL → restore → assert PASS) and the working tree is verifiably clean BEFORE any unrelated commit may proceed in the same checkout.
3. Concurrent subagents working in the SAME checkout (single working tree, no `git worktree add`) MUST coordinate through a lockfile or marker file (e.g. `.git/MUTATION_IN_PROGRESS`) so a logo-fix subagent cannot race a JWT-mutation subagent. The constitution submodule's `scripts/quiescence_check.sh` (when implemented) is the universal lock helper.
4. When parallel work is necessary, the §11.4.20 / §11.4.70 subagent-driven mandate SHOULD be combined with `git worktree add` to give each subagent an isolated working tree — eliminates the cross-mutation race by construction.
5. Post-commit, an automated `mutation-residue-scanner` (consuming project's `scripts/mutation_residue_audit.sh`) MUST run before push. Any commit containing a mutation marker → push BLOCKED, commit MUST be reverted or amended before mirrors are updated.

Composes with §1.1 (mutation-paired gates — quiescence rule protects the mutation cycle from concurrent contamination), §11.4.20 / §11.4.70 (subagent-driven-by-default — quiescence rule is what makes parallel subagent dispatch safe), §11.4.27 (no-fakes-beyond-unit — a mutation residue swept into a commit IS a fake-pass surface in production), §11.4.10 (credentials handling — same class of "do not let unrelated content leak into a commit"), §107 (a security-bypass mutation

that ships to production is the gravest §107 PASS-bluff: tests pass against the mutated codebase, but the end user loses an entire security property), §11.4.71 (pre-push fetch + investigate — quiescence check is the pre-push companion).

Classification: universal (per §11.4.17). Every project under this Constitution carries this rule.

Canonical authority: constitution submodule `Constitution.md` §11.4.84.

Non-compliance is a release blocker. No `--allow-residue`, `--skip-quiescence`, `--mutation-cleanup-later` flag exists. A mutation marker that lands in a tagged commit is a critical defect regardless of how briefly it persisted.

2026-05-24

§ . . — Stress + Chaos Test Mandate (User mandate, — —)

Short tag: `stress-chaos-mandate` .

Forensic anchor (verbatim user mandate, 2026-05-24):

"Every fix or improvement you do MUST BE covered with full automation stress and chaos tests so we are sure nothing can break the functionality and all edge cases are monitored and polished and additionally fixed if that is needed! Everything must produce rock solid proofs and follow fully no-bluff policy!"

The mandate. Every fix or improvement landed in a consuming project MUST ship with full-automation **stress** AND **chaos** test suites that exercise edge cases, sustained load, concurrent contention, and failure-injection. A fix that PASSES its happy-path test but has never been exercised under stress or under fault-injection is a §11.4 / §107 PASS-bluff at the resilience layer: it claims to work but carries no evidence that real-world adversarial conditions (sustained throughput, parallel contention, partial failure, resource exhaustion, malformed input) leave the fix intact and the user-visible behaviour correct.

Definitions (closed-set, mechanically auditable):

1. **Stress test** — exercises the fix-under-test under sustained or concurrent load above ordinary usage. At MINIMUM one of:
 - **Sustained load** — $N \geq 100$ sequential iterations OR ≥ 30 seconds wall-clock continuous load. Per-iteration latency MUST be recorded; percentile distribution (p50/p95/p99) MUST be reported.
 - **Concurrent contention** — $N \geq 10$ parallel invocations. All N MUST complete. No deadlock, no resource leak (file-descriptor count, process-table count, RSS), no data race in shared state.
 - **Boundary conditions** — minimum input (empty), maximum input (provider/codec/protocol-limit-bound), boundary input (off-by-one at every size threshold). Each boundary MUST produce a categorised result (success, categorised-error, deterministic-skip) — NEVER an uncaught exception or silent corruption.
2. **Chaos test** — exercises the fix-under-test under failure-injection per the §11.4.69 closed-set evidence taxonomy. At MINIMUM one chaos category appropriate to the fix-class:
 - **Process-death injection** — kill the primary process, dependency process, or upstream service mid-operation. Recovery path MUST be deterministic + categorised (clean error, automatic restart, circuit-breaker open).
 - **Network-fault injection** — drop / delay / reorder packets on the relevant interface mid-call. Categorised error per §11.4.69 (e.g. `category=network` / `category=upstream`).
 - **Input-corruption injection** — corrupt the input file / config / .env mid-test. Test MUST detect + report the corruption — NEVER silently consume corrupted input.
 - **Resource-exhaustion injection** — fill disk to 99 %, allocate memory pressure (OOM-injection), exhaust file descriptors, exhaust sockets. Fix MUST refuse new work cleanly OR degrade gracefully — NEVER crash with uncaught error.
 - **State-corruption injection** — mid-flight database lock loss, mid-flight file-system partial-write error, mid-flight cache invalidation. Recovery path MUST restore consistent state.

Anti-bluff (mandatory):

1. Every stress + chaos test PASS MUST cite a captured-evidence artefact path per §11.4.5 + §11.4.69. Acceptable evidence: `per-iteration latency.json` / `throughput.csv` / `categorised_errors.txt` / `state_delta_snapshot.json` / `process_lifecycle.log` / `recovery_trace.log` . Metadata-only PASS ("all iterations exited 0") without per-iteration latency / failure / state evidence is itself a §11.4 PASS-bluff at the stress-layer.
2. Helper library: every consuming project SHOULD ship a `stress_chaos.sh` (or analogous language-binding library) exposing reusable primitives — `ab_stress_run` , `ab_stress_concurrent` , `ab_chaos_kill_pid_during` , `ab_chaos_drop_network_during` , `ab_chaos_corrupt_file_during` , `ab_chaos_oom_pressure_during` , `ab_chaos_disk_full_during` . Each helper composes with `ab_pass_with_evidence` / `ab_skip_with_reason` per §11.4.69.
3. Chaos-injection cleanup is non-negotiable. A test that corrupts a `.env` MUST restore it in `trap '...' EXIT` . A test that fills the disk MUST `rm` the filler in EXIT. A test that kills a process MUST verify the process is restarted (or explicitly leave it down with operator-visible reason). Cleanup failure = §11.4.14 (test-playback cleanup) violation.

4-layer coverage per §11.4.4(b):

- **Pre-build gate** — stress + chaos test files exist + executable + parseable under `sh -n` + `bash -n` per §11.4.67; helper library exists + executable; the fix's pre-build gate cites the stress + chaos test file path.
- **Paired meta-test mutation** — per §1.1, removing the chaos-injection step or stripping the per-iteration evidence-capture → the gate FAILs.
- **On-device test** (if the fix is LIVE_ADB_TESTABLE per §11.4.51) — the stress + chaos test is dispatched against a real device, captured-evidence directory under `qa-results/<run-id>/stress_chaos/` .
- **HelixQA Challenge entry** (if the fix is a user-visible feature per §11.4.4(b) layer 4) — a Challenge bank entry references the same stress + chaos test as the validation route.

Composition with related §11.4 anchors:

- **§11.4 / §107** — stress + chaos test PASS is the strongest end-user-quality signal: it demonstrates the fix survives adversarial conditions the happy-path test cannot reach.
- **§11.4.1** — FAIL-bluffs forbidden. A stress test that crashes with `set -u` because the chaos-injection step left a variable unset is a FAIL-bluff; fix at source.
- **§11.4.5** — captured-evidence content quality applies to stress + chaos evidence too (latency distribution recorded, not just count; error categories enumerated, not just "some errors").
- **§11.4.6** — no guessing. Categorised errors per closed-set. "Probably the network" is forbidden — capture the categorisation.
- **§11.4.43** — TDD RED-first. The stress + chaos suite SHOULD start RED: write the test, observe the failure under load/chaos, land the fix, observe the GREEN.
- **§11.4.50** — deterministic consistency. The stress test's N iterations MUST produce identical exit codes AND identical evidence-hashes (volatile prefixes stripped) per §11.4.50.
- **§11.4.52** — autonomous validation. Stress + chaos tests MUST run end-to-end without operator presence.
- **§11.4.69** — universal sink-side positive-evidence taxonomy. Every stress + chaos PASS uses `ab_pass_with_evidence` with a real captured-evidence file path.
- **§11.4.83** — `docs/qa/` end-user transcript discipline composes with chaos-test evidence (chaos-test recovery transcripts ARE end-user-channel proofs).

Classification: universal (per §11.4.17). Every consuming project carries this rule. Stress + chaos discipline is language-/platform-/build-system-agnostic.

Canonical authority: constitution submodule `Constitution.md` §11.4.85.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--skip-stress`, `--no-chaos`, `--happy-path-suffices`, `--stress-test-later` flag exists. The discipline exists because the User mandate is unambiguous: "Everything must produce rock solid proofs and follow fully no-bluff policy!" — and a fix never tested under stress + chaos is not a rock-solid proof.

§ 11.4.45 – Roster/corpus-backed Status-doc auto-sync mandate (User mandate, – –)

Forensic anchor — verbatim user mandate (2026-05-25):

"Make sure that assets and players Status docs are ALWAYS regularly updated and in sync like all others Status docs — any time we add or modify the assets content(s) or we change or add new / remove existing pre-installed video and audio player apps! This MUST WORK OUT OF THE BOX!"

Some Status docs (§11.4.45) are backed not by hand-written narrative alone but by a **tracked roster** (a set of installed apps / components) or a **tracked asset corpus** (a directory of test / media assets). For these, freshness cannot be left to operator vigilance — the moment a member of the roster/corpus changes (a player app added / removed / renamed; a test asset added / modified / removed) the Status doc, its Status_Summary, and their HTML + PDF exports MUST be brought back into sync **out of the box**, mechanically, not by remembering to.

The mandate. Every Status doc whose subject is a tracked roster or asset corpus MUST be kept in sync with that subject by a mechanism where **all** of the following hold:

1. **Drift-proof signal.** A content *fingerprint* — sha256 of the sorted member list, NOT mtime (which `git checkout` resets) — is persisted in a sidecar next to the Status doc. Adding / removing / modifying a member changes the fingerprint deterministically.
2. **Sync helper.** A helper regenerates the fingerprint and re-exports HTML + PDF (composing with the §11.4.65 universal exporter). Invocable standalone AND wired so the sync happens automatically (commit path and/or pre-build).
3. **Enforcement gate.** A pre-build gate recomputes the live fingerprint and FAILs when it differs from the persisted one — forcing the Status doc to be updated whenever the roster/corpus changes. Mirrors §11.4.12 `CM-ISSUES-SUMMARY-SYNC` + §11.4.45 `sync_integration_status` .

4. **Paired §1.1 mutation.** A meta-test mutation corrupts the fingerprint (introducing drift without a doc update) and asserts the gate FAILs — proving the gate is not itself a bluff.

Composes with §11.4.12 (auto-generated docs sync — the canonical sibling pattern), §11.4.45 (integration-status-doc — §11.4.86 is its roster/corpus specialisation), §11.4.53 + §11.4.56 (Summary parity), §11.4.57 + §11.4.59 (README doc-link freshness), §11.4.60 (composite doc sync), §11.4.65 (universal Markdown export), §11.4.6 (no-guessing — the fingerprint is FACT, not a guess about freshness).

Classification: universal (per §11.4.17) — the principle (roster/corpus-backed Status docs auto-sync on subject change) is project-agnostic; the consuming project supplies the specific docs, roster/corpus sources, helper, and gate name in its own CLAUDE.md / AGENTS.md / QWEN.md per §11.4.35.

Canonical authority: constitution submodule `Constitution.md` §11.4.86.

Non-compliance is a release blocker regardless of context. No escape hatch —
no `11--skip-roster-sync4, 87--allow-status-drift, --roster-sync-not-applicable` flag exists.

2026 05 26

§ . . — Endless-loop autonomous work + zero-idle agent dispatch + anti-bluff testing mandate (User mandate, — —)

Forensic anchor — verbatim User mandate (2026-05-26):

"when we say to you that all work MUST BE continued in endless loop until there is no any open items, no unfinished workable items from our Issues docs, or from Continuation document or any unfinished work by agents at the moment it means that you will ALWAYS run until there is nothing workable left to be done in endless loop, fully autonomously. You will spawn agents or agents-driven work whenever that is possible or required! Not a single agent or main work stream will sit idle except if it waits for the results of something - some execution and similar! All work MUST BE always covered with comprehensive tests - all supported tests for the project which produce real proofs for the contexts and functionalities they are testing! Bluff of any kind is not allowed and all work and tests MUST WORK completely as proofs (physical) driven and in complete anti-bluff manner!"

When an operator instructs an AI agent to "continue in endless loop fully autonomously" — or any semantically-equivalent phrasing — the agent MUST interpret this as a HARD-CONTRACT execution covenant covering five non-negotiable obligations:

(A) Endless-loop continuation. The agent MUST continue working in the autonomous-loop framework until ALL of the following are simultaneously TRUE:

- docs/Issues.md Status-column has zero In progress / Ready for testing / In testing / Reopened entries (§11.4.15 closed-set values).
- docs/CONTINUATION.md §3 "Active work" section is empty.
- No background subagent is mid-execution (TaskList reports all tasks completed or pending with pending items NOT being prerequisite-blockers for anything actionable).
- No external dependency is in-flight (build, flash, push, sync).

If any of those conditions is FALSE, the agent MUST continue working — by claiming the next item from the priority queue, by dispatching a subagent for parallelisable work, or by polling a background task that will be the wake signal. An agent that schedules a wake-up and ends the turn while Issues.md still has actionable items, OR while CONTINUATION.md §3 still has active work, OR while a subagent is still working productively, is in violation.

(B) Zero-idle agent dispatch. When the agent identifies parallelisable work that does not contend on the same file-scope as the main work stream, the agent **MUST** dispatch a background subagent rather than serialising the work. The main work stream + every background subagent operate concurrently. "Wait for the results of something" is the **ONLY** acceptable reason for an agent (main or subagent) to be idle. The §11.4.20 + §11.4.58 + §11.4.70 subagent-driven covenants compose with this anchor.

(C) Comprehensive test coverage with real (physical) proofs. Every workable item closed in the loop **MUST** land four-layer test coverage per §11.4.4(b) — pre-build gate + on-device test + paired §1.1 meta-test mutation + HelixQA Challenge bank entry — and every test PASS **MUST** cite a captured-evidence artefact per §11.4.5 + §11.4.69. The phrase "physical proofs" is interpreted as: captured audio (`tinycap` WAV with RMS + `ffprobe` channels), captured video (screen recording + `ffprobe` frames + recording-analyzer event match), captured network state (`dumpsys` + sink-side probe per §11.4.13), captured UI state (`uiautomator` dump per §11.4.48), captured sysfs state (`/sys/...` / `/proc/...` snapshots). Metadata-only PASS, configuration-only PASS, absence-of-error PASS, grep-without-runtime PASS are all critical defects regardless of the green summary line.

(D) Anti-bluff guarantee end-to-end. The §11.4 covenant family (§11.4.1 FAIL-bluffs / §11.4.2 recorded-evidence / §11.4.6 no-guessing / §11.4.7 demotion-evidence / §11.4.27 no-fakes-beyond-unit / §11.4.50 deterministic-consistency / §11.4.52 autonomous-validation / §11.4.68 audio sink-side / §11.4.69 universal sink-side / §11.4.83 docs/qa/ end-user evidence) is the operative truth-discipline of the endless-loop covenant. Every closure narrative **MUST** cite captured-evidence; every demotion **MUST** cite same-conditions retest; every status transition **MUST** cite the deterministic-consistency baseline. **Tests AND Challenges (HelixQA) are bound equally** — a Challenge that scores PASS on a non-functional feature is the same class of defect as a unit test that does. The User mandate explicitly invokes the historical forensic anchor: "we had been in position that all tests do execute with success and all Challenges as well, but in reality the most of the features does not work and can't be used! This **MUST NOT** be the case and execution of tests and Challenges **MUST** guarantee the quality, the completion and full usability by end users of the product!"

(E) Termination criteria. The endless-loop terminates ONLY when:

1. ALL conditions in (A) are simultaneously TRUE (zero open items + zero active CONTINUATION work + zero subagents in flight + zero external dependencies in-flight); OR
2. The operator explicitly issues a `STOP / END LOOP / pause autonomous work` instruction; OR
3. The host-session-safety covenant (§12 + §12.6 + §12.7 + §12.8 + §12.9 + §12.10) demands suspension to protect the operator's session; OR
4. The agent has scheduled a wake-up to resume against a known-future-actionable signal (e.g. CI completion, build completion, hardware-attended phase).

No escape hatch. No `--idle-OK` , `--skip-endless-loop` , `--bluff-permitted-for-this-task` , `--metadata-only-test-suffices` , `--no-physical-proof-required` flag exists. The covenant is uniform across every iteration.

Composes with §11.4 (anti-bluff covenant — historical forensic anchor of the project), §11.4.1 (FAIL-bluffs), §11.4.2 (recorded-evidence), §11.4.4 (four-layer coverage), §11.4.5 (audio + video 5-layer quality), §11.4.6 (no-guessing), §11.4.7 (demotion-evidence), §11.4.20 + §11.4.58 + §11.4.70 (subagent-driven defaults), §11.4.27 (no-fakes-beyond-unit + 100% test-type coverage), §11.4.42 (iteration-discipline), §11.4.43 (TDD-fix), §11.4.50 (deterministic-consistency), §11.4.52 (autonomous-validation), §11.4.68 (audio sink-side), §11.4.69 (universal sink-side), §11.4.83 (docs/qa/ end-user evidence), §11.4.85 (stress + chaos), §11.4.86 (roster/ corpus auto-sync), §12 family (host-session safety boundaries), §12.10 (CONTINUATION-document maintenance — the source-of-truth state the loop checks against).

Pre-build gate `CM-COVENANT-114-87-PROPAGATION` enforces this anchor literal in every CLAUDE.md / AGENTS.md / QWEN.md across the canonical fleet (parent + consumer submodules). Paired meta-test mutation strips the literal `11.4.87` → gate FAILs.

Canonical authority: this Constitution.md §11.4.87 in the HelixConstitution submodule (`git@github.com:HelixDevelopment/HelixConstitution.git`) — inherited per §11.4.35 by every consuming project's repo-root CLAUDE.md / AGENTS.md / QWEN.md.

Non-compliance is a release blocker regardless of context. Skipping the loop, idling without a waited-on signal, or accepting bluff-evidence PASS is severity-equivalent to a §11.4 PASS-bluff at the project-execution layer.

2026 05 26

§ . . — Background-push mandate:
commit-lock release immediately after commit,
push runs detached (User mandate,
— —)

Forensic anchor — verbatim User mandate (2026-05-26):

"Make sure all these pushes are being done ALWAYS in backgroubd in parallel with main work stream so we do not loose time waiting. Everything is committed anyway? If we can do this add this as mandatory rule / constraint that we must respect and follow ALWAYS! ... We MUST ensure that main work stream has always something to do, or wait for the results only when that is absolutely required! We have enormous amount of work in front of us, short deadlines and we MUST DELIVER!"

Forensic incident. 2026-05-26 between 14:35Z and ≥19:30Z (~5 hours observed), a single `commit_all.sh` invocation held its `.git/.commit_all.lock` flock the entire time — the COMMIT had succeeded within seconds at `63d307234ce` , but the synchronous `do_push` call kept the flock while `git push vasicdigital_gitlab` did a large initial-mirror upload. **Every subsequent `commit_all.sh` invocation in that window was BLOCKED** — 21 submodule pointer bumps + parent §11.4.87 propagation work + §KA Critical audio quality matrix artifacts all sat staged in the working tree waiting for the push to finish. This is the exact anti-pattern §11.4.87 (zero-idle) prohibits at the project-execution layer.

The mandate.

(A) Flock release IMMEDIATELY after commit lands. Once `git commit` returns 0 (commit object recorded in the local repo), the `.git/.commit_all.lock` MUST be released BEFORE `do_push` is invoked. The commit is durable on local disk regardless of push outcome — there is no correctness reason to gate further local work on a remote-push round-trip.

(B) Push runs detached. After flock release, the push step is spawned via `nohup ./push_all.sh ... > <log> 2>&1 &` then `disown` so the orchestrator's exit does not propagate SIGTERM. The orchestrator's exit code reports COMMIT success, NOT push success (push success is logged separately and reconciled by §11.4.71 fetch-before-edit on the next interactive turn).

(C) Push concurrency control: per-remote serialization, multi-remote parallelism. `push_all.sh` MUST acquire a per-remote flock (`.git/.push.<remote>.lock`) so two concurrent invocations targeting the same remote serialize (avoid push-race / non-fast-forward), but invocations targeting DIFFERENT remotes run in parallel. This converts the prior 5-hour serial-stall pattern into a multi-stream pipeline where the operator's main work stream never blocks on the slowest mirror.

(D) Failure surface. Backgrounded push failures land in `qa-results/push_failures/<timestamp>_<remote>.log`. The next interactive autonomous-loop tick MUST check that directory (per §11.4.87(A) "no external dependency is in-flight" check) and surface failures — re-push attempted automatically, or operator notified if auth/quota issue. Silent push-failure is a §11.4 PASS-bluff at the distribution layer.

(E) Synchronous-push escape. An explicit `--sync-push` CLI flag on `commit_all.sh` preserves the legacy synchronous behaviour for `--force-with-lease` paths (per §11.4.41 force-push merge-first audit) where the push outcome MUST gate the next step. This is the ONLY escape — without it, every push is backgrounded.

Composes with §2.1 (multi-upstream push norm — §11.4.88 makes it non-blocking), §9.2 (data safety — backgrounded push still uses hardlinked-backup pre-op), §11.4.41 (force-push merge-first — uses `--sync-push` escape), §11.4.42 (iteration discipline — endless-loop ticks no longer waste budget on remote-network latency), §11.4.71 (fetch-before-push merge-first — applies before

background push spawns), §11.4.87 (endless-loop / zero-idle — §11.4.88 is the implementation seam that makes §11.4.87(B) genuinely zero-idle for the dominant blocker class).

Pre-build gate `CM-COVENANT-114-88-PROPAGATION` enforces this anchor literal in every `CLAUDE.md` / `AGENTS.md` / `QWEN.md` across the canonical fleet. Pre-build gate `CM-BACKGROUND-PUSH-WIRED` verifies `commit_all.sh` releases flock before `do_push` AND spawns push via `nohup ... & AND push_all.sh` acquires per-remote flock. Paired §1.1 meta-test mutations strip the load-bearing literals → gates FAIL.

Canonical authority: this Constitution.md §11.4.88 in the HelixConstitution submodule (`git@github.com:HelixDevelopment/HelixConstitution.git`).

Non-compliance is a release blocker. Synchronous push (without `--sync-push` flag) is severity-equivalent to a §11.4 PASS-bluff at the project-execution layer. No escape hatch beyond `--sync-push` for §11.4.41 force-push events.

2026 05 27

§ . . — Background test execution mandate (User mandate, — —)

Forensic anchor — verbatim user mandate (2026-05-27):

"Any tests we are executing, especially long test cycles, MUST BE performed in background in parallel with main work stream! This MUST NOT block our capabilities to work on queued workable items (and tackle main Issues document(s)). Make sure we strictly follow this starting now! Main work stream can be blocked or sit iddle only if absolutely needed and if it depends hard on results of some background execution."

Forensic incident. 2026-05-27 the conductor invoked `bash device/rockchip/rk3588/tests/pre_build_verification.sh` synchronously with a `timeout 360` foreground wrapper — even though the operator had previously mandated §11.4.87 endless-loop zero-idle. The 6-7 minute test run blocked the entire main work stream

from progressing on §JV/§JW/§JX/§JY scaffolding + further Issues.md drains. This is the exact anti-pattern §11.4.87(B) prohibits at the test-execution layer (analogous to §11.4.88 at the push-execution layer).

The mandate.

(A) Long-running tests run detached. Any test cycle expected to exceed ~30 seconds — `pre_build_verification.sh` , `meta_test_false_positive_proof.sh` , `test_all_fixes.sh` , `recent_work_validate.sh` , HelixQA Challenge banks, on-device 4-phase cycles, full-suite retests per §11.4.40, `audio_loop_supervisor.sh`, `dual_display_record.sh` harnesses, `verify-all-constitution-rules.sh` — MUST be spawned via `nohup ... > <log> 2>&1 & + disown` and `<log>` placed under a known directory (e.g. `qa-results/<test_id>_<timestamp>.log`) for later inspection.

(B) Main stream proceeds on queued items. The conductor MUST immediately return to the §11.4.42 priority queue (open Issues.md items, queued PWUs, doc-sync work, audit work, anything not depending hard on the in-flight test's exit code). "Wait for results" is the ONLY acceptable idle reason per §11.4.87(B) — and even then, the conductor checks back on the test result via short polling (`tail / mtime / exit-status` file) rather than blocking on `wait` .

(C) Hard-dependency gating. A subsequent step that genuinely requires the test's exit code (e.g. `commit_all.sh` refuses if meta-test is in flight per §11.4.84; release-tag gate per §11.4.40 step 4) MUST express that dependency explicitly — poll the exit-status file or check `pgrep -af <test-name>` before proceeding. If the test is still running and the next step needs its result, surface that to the operator per §11.4.66 with interactive options.

(D) Failure surface + cleanup. Backgrounded test failures land in their `<log>` files. The next autonomous-loop tick MUST check that directory (analogue of §11.4.88(D)). Test-script-level cleanup (per §11.4.14 + §11.4.84 + this anchor's restore discipline) is non-negotiable — `trap '<cleanup>' EXIT INT TERM` in every long-running test that mutates global state (mid-mutation backup files, sysfs nodes, ALSA states, etc.).

(E) Pre-empt the blocking pattern. The conductor MUST NOT invoke `bash <long-test>.sh` synchronously when a backgrounded variant suffices. Foreground execution is permitted ONLY when: (a) the test completes in <30 s, OR (b) the operator explicitly requests foreground execution for a specific step.

Synchronous execution of a multi-minute test absent operator authorisation is a §11.4.89 violation regardless of how much "context-cache benefit" the foreground path appears to offer.

(F) Concurrency control. Multiple backgrounded tests of the SAME script MUST serialise via per-script flock (e.g.

`.git/.pre_build.lock` , `.git/.meta_test.lock`). Concurrent invocations of DIFFERENT tests run in parallel.

Composes with §11.4.42 (iteration discipline — backgrounding tests recovers the operator's wait-time per §11.4.82(F) parallel multi-device testing), §11.4.66 (interactive clarification — surface hard-dependencies as options instead of blocking), §11.4.82 (iteration-speedup discipline — §11.4.89 is the dominant blocker class at the test layer, analogous to §11.4.88 at the push layer), §11.4.84 (working-tree quiescence — backgrounded mutation tests still need quiescence checks), §11.4.85 (stress + chaos — long stress runs are quintessential §11.4.89 candidates), §11.4.87 (endless-loop zero-idle — §11.4.89 is the implementation seam at the test-execution layer that makes §11.4.87(B) genuinely zero-idle for the test blocker class), §11.4.88 (background-push — §11.4.89 is the symmetric anchor for tests).

Pre-build gate `CM-COVENANT-114-89-PROPAGATION` enforces this anchor literal across the canonical fleet. Pre-build gate

`CM-BACKGROUND-TEST-EXECUTION-WIRED` verifies that long-running test invocations in helper scripts use `nohup ... & + disown` pattern (or equivalent). Paired §1.1 meta-test mutations strip the load-bearing literals → gates FAIL.

(G) Subagent long-op discipline + conductor-owns-inherently-long-ops (extension, 2026-07-15). Forensic FACT (2026-07-15): background subagents were repeatedly KILLED by the runtime's no-progress stream watchdog (~600 s with no tool-call output) because each had run ONE synchronous long operation inside a single tool call (a full pre-build sweep / a long test suite / a device flash / a hung network call) — the host and the network were ruled out as causes; the watchdog fired on the SILENCE, not on any failure. Therefore: (1) a SUBAGENT MUST NOT run any single synchronous operation expected to exceed the runtime's no-progress watchdog budget (conservatively ~5 min) — it MUST background the operation per clause (A) (`nohup ... & + disown`) and POLL with short, frequent tool calls (each poll is the progress the watchdog observes); (2) inherently-long operations that a subagent CANNOT usefully background-and-poll (full artifact

builds, full test suites, device flashes / deploys) are the CONDUCTOR's to run — the conductor (main session) carries no such watchdog and MUST own them, dispatching subagents only for the parallelisable slices around them; (3) a subagent killed by the stall watchdog is a CRASH per §11.4.147 — its work is NOT lost, NOT done, NOT forgotten: the registry entry flips to `crashed` and the unit is RESPAWNED until `complete` ; a stall-kill silently absorbed as a completion is a §11.4.147 + §11.4 bluff. Honest boundary (§11.4.6): the watchdog budget is a RUNTIME property — read it from the runtime, never guess it; when it is unknown, treat any operation > ~5 min as long.

Canonical authority: this Constitution.md §11.4.89 in the HelixConstitution submodule (`git@github.com:HelixDevelopment/HelixConstitution.git`).

Non-compliance is a release blocker. Synchronous long-test execution (without operator authorisation) is severity-equivalent to a §11.4 PASS-bluff at the project-execution layer. No escape hatch beyond explicit per-invocation operator authorisation.

2026 05 27

§ . . — Obsolete status + per-item
obsolescence audit mandate (User mandate,
- -)

Forensic anchor — verbatim user mandate (2026-05-27):

"Regarding the Critical Bug No 6 - Albums cover etc. - Now it seems obsolete after latest request for new behavior when audio and video content is being played. If we do not have Obsolete as an option / status for the ticket in its resolving, we should introduce this now and mark obsolete tickets with some light gray background or something that will indicate that item is no longer valid. Maybe text - the description to be strikethrough styled as well! Once all this is done please review all existing open or resolved workable items (fixed as well) if they are obsolete - not valid any more. ... There MUST NOT be any mistake! No bluff is allowed of any kind!"

The §11.4.15 Status closed-set is extended with a 4th terminal value `obsolete` (`→ Fixed.md`) matching the existing §11.4.33 closure-vocabulary table (`Bug → Fixed`, `Feature → Implemented`, `Task → Completed`). The `obsolete` value applies regardless of Type when the workable item is **no longer valid** because: (a) the underlying behaviour the item described has been superseded by a later operator mandate / design pivot (canonical example: Bug #6 Albums cover, superseded by §JU + the 2026-05-27 secondary-display video-playback redesign), (b) the affected feature was REMOVED from the product, (c) the item was duplicated and the canonical entry resolved it, (d) the item was filed against a hardware/topology variant we no longer support.

Mandatory obsolescence audit metadata — every `obsolete` (`→ Fixed.md`) heading MUST carry, within 8 non-blank lines of the heading, an `**obsolete-Details:**` line capturing four sub-facts (mirroring §11.4.21 + §11.4.34 audit-line patterns):

- **Since:** ISO date (`YYYY-MM-DD`).
- **Reason:** one-line cause from the closed vocabulary
`{ superseded-by-design-change | superseded-by-later-mandate | feature-removed | duplicate-of | unsupported-topology | not-reproducible } . (not-reproducible = a reported defect that does NOT reproduce on the canonical tree / baseline — an environment / isolated-worktree artifact (PATH / shell / missing-checkout), not a real product defect; the triple-check evidence MUST capture the canonical-tree non-reproduction, per §11.4.6 no-guessing + §11.4.7 demotion-evidence.)`
- **Superseding-item:** §-letter reference or User-mandate verbatim quote anchor citing the work that obsoleted it.
- **Triple-check evidence:** path to captured evidence (git log, code grep, runtime behaviour) confirming the item is genuinely no longer valid. Per operator mandate "There MUST NOT be any mistake" — bare assertion is forbidden; positive captured evidence per §11.4.6 is mandatory.

Visual treatment — the §11.4.23 colorizer MUST style `obsolete` cells with `cell-status-obsolete` class: light-gray background `#E0E0E0` + strikethrough text on the row's description cell. Operator mandate verbatim: "mark obsolete tickets with some light gray background or something that will indicate that item is no longer valid. Maybe text - the description to be strikethrough styled as well".

Audit cadence — at every release-gate sweep (per §11.4.40), the conductor MUST re-evaluate every non-terminal Issues.md item AND every Fixed.md item for obsolescence per the closed-set Reason vocabulary above. New obsolescence findings produce `Obsolete` ($\rightarrow$ Fixed.md) migrations atomic per §11.4.19. Triple-check is non-negotiable: each obsolescence claim cites (i) the superseding mandate, (ii) code-state grep confirming the obsoleted behaviour is gone, (iii) runtime/test evidence confirming no test still exercises the obsoleted path.

Pre-build gates `CM-COVENANT-114-90-PROPAGATION` (anchor literal across canonical files) + `CM-ITEM-OBSOLETE-DETAILS` (every `Obsolete` heading carries `**Obsolete-Details:**` line within 8 lines) + `CM-OBSOLETE-COLORIZER-WIRED` (§11.4.23 colorizer emits `cell-status-obsolete` class on the right cells + the CSS file defines the class). Paired §1.1 mutations strip the anchor literal / delete an Obsolete-Details line / strip the colorizer class — every mutation FAILs its gate.

Composes with §11.4.15 (item-status closed-set extension), §11.4.16 (Type tracking unchanged — Obsolete is orthogonal to Type), §11.4.19 (Fixed-document column-alignment — Obsolete migrates to Fixed.md just like Fixed/Implemented/Completed), §11.4.21 (Operator-blocked details-line pattern — Obsolete-Details mirrors), §11.4.23 (visual cue + grouping — adds `cell-status-obsolete`), §11.4.33 (closure vocabulary — Obsolete is the 4th terminal value), §11.4.34 (Reopened-source pattern — Obsolete uses analogous audit line), §11.4.40 (release-gate sweep — audit cadence trigger), §11.4.42 (iteration discipline — obsolescence audit is a routine cycle step), §11.4.66 (interactive clarification — surface ambiguous obsolescence to operator), §11.4.71 (pre-push fetch — re-check obsolescence post-fetch).

Canonical authority: this Constitution.md §11.4.90 in the HelixConstitution submodule.

Non-compliance is a release blocker. Marking a still-valid item Obsolete is a §11.4 PASS-bluff at the planning layer. Marking a genuinely-obsolete item with a non-terminal Status is documentation drift. No escape hatch — every obsolescence finding follows the triple-check evidence rule above.

§ . . – Summary-doc clarity mandate (User mandate, – –)

Forensic anchor — verbatim user mandate (2026-05-27):

"We see in Summary docs - Issues_Summary some not clear one line descriptions - like for example: 'Composes with'. For each workable item in any variant of documentation - long or short descriptions - we MUST HAVE clearly understandable meaning of particular entry of the document! Re-evaluate all of them, fix the descriptions (especially one liners) and make sure that every team member can clearly understand what that particular or any particular workable item from the list is exactly about! There cannot be misunderstanding or unclarity of any kind and no bluff allowed!"

Every short-form summary entry (Issues_Summary.md, Fixed_Summary.md, README.md doc-link section, Status_Summary.md page 1 + 2, every other one-liner derived from a long-form tracker entry) MUST contain a description that is **self-contained + meaningfully informative** about WHAT the item is, not a fragment of the long-form body. Specifically forbidden one-liner anti-patterns:

- **Section labels** as descriptions: Composes with , Closure criteria , Fix direction , Forensic anchor , Reopened-Details , Operator-Block-Details , Obsolete-Details , Summary , Status , Type , Severity , etc. — these are section *labels* inside the long-form entry, NEVER appropriate as the one-liner description.
- **Bare numeric / categorical fragments:** Critical , High , Medium , Low , Bug , Feature , Task , Fixed , standalone — these are metadata column values, not descriptions.
- **Generic restatements of Status:** In progress , Queued , Reopened , Operator-blocked , Fixed , standalone — again metadata column values.
- **Section-marker echoes:** anything matching `^[A-Z][a-z]+ (with|by|on|in|to|for):?$` patterns where the right-hand side is empty.
- **§-letter alone:** §KB , §EU , etc. — the section letter is the identifier, not the description.

Required one-liner shape — every summary entry's description column **MUST** contain a complete clause (≥ 6 words OR ≥ 40 characters, whichever is longer) that names the SUBJECT (component / app / behaviour / defect class) + the PROBLEM or GOAL (what's broken / what's being added / what's being audited). Example transformations:

- Bad: "Composes with"
- Good: "MPV fork comprehensive secondary-display playback + interaction-pattern + chaos test suite (§KB-2)"
- Bad: "Critical"
- Good: "HDMI audio total loss with 'command error 1' on D3 post-reflash (Tier 1/2/3 source-side fixes landed, REQUIRES_REBUILD validation pending)"
- Bad: "\$JU"
- Good: "Presenter 2nd-display behaviour simplification + wake-on-play + notification dim toggle (umbrella, REQUIRES_REBUILD pending validation)"

Generator-level enforcement — `generate_issues_summary.sh` + `generate_fixed_summary.sh` + `update_readme_doc_links.sh` + `generate_status_summary.sh` **MUST** extract the one-liner from the **H1/H2 heading line** of the source long-form entry (which already encodes the subject + context per §11.4.15 + §11.4.16 + §11.4.54 ATM-NNN convention), **NEVER** from arbitrary downstream text. The generators **MUST** refuse to emit a row whose description matches any of the forbidden anti-patterns above — emitting `(MISSING DESCRIPTION – fix source heading)` placeholder instead, with the offending row visually highlighted.

Pre-build gate `CM-SUMMARY-CLARITY-DESCRIPTIONS` scans every `*_Summary.md` for the forbidden anti-patterns above; any match **FAILS** the gate. Paired §1.1 mutation injects "Composes with" into a Summary cell → gate **FAILS**.

Audit cadence — at every release-gate sweep (§11.4.40 + §11.4.42 step 4), re-evaluate every summary row for clarity. Operator-facing exports (HTML + PDF) **MUST** never ship with an anti-pattern row.

Composes with §11.4.12 (Issues_Summary sync — clarity is a property of the generated row), §11.4.19 (Fixed-document column alignment — clarity applies symmetrically), §11.4.23 (colorizer — anti-pattern rows highlighted), §11.4.44

(revision header — generator updates revision on clarity regeneration), §11.4.53 (Fixed_Summary parity), §11.4.56 (Status_Summary two-audience format — page 1 audience MUST get the clearest descriptions), §11.4.57 (README doc-link rows — same anti-pattern rule), §11.4.59 (README always-sync), §11.4.60 (composite docs always-sync), §11.4.65 (universal MD export), §11.4.74 (mechanical enforcement — pre-commit hook for summary clarity).

Canonical authority: this Constitution.md §11.4.91 in the HelixConstitution submodule.

Non-compliance is a release blocker. A summary row whose description is an anti-pattern fragment is severity-equivalent to a §11.4 PASS-bluff at the documentation layer — it implies the item is documented when it is not. No escape hatch. 92

2026 05 27

§ . . — Multi-pass change-evaluation discipline (User mandate, — —)

Forensic anchor — verbatim user mandate (2026-05-27):

"Every change to the project or codebase we do MUST BE evaluated in several passes and in in-depth analysis for potential new issues or problems it can introduce! We MUST BE sure that: change nothing breaks, we do validate and verify this in codebase and through comprehensive deep research, that no new issues or problems of any kind is bringing and that main task is achieved with particular change with no bluff of any kind! After we do change or set of changes this mandatory steps MUST BE taken! Then, and only then, when we have validate and verified codebase and performed full analysis and taken into the account all existing knowledge(s) and information sources and confirmed all points we have mentioned now, then this can be accepted as change that is going to be committed, pushed, tested, released and so on! Write as many as required additional documentation and notes about our codebase, the system, or the changes so we can achieve this goals!"

Every non-trivial change to the project codebase MUST pass a multi-pass evaluation BEFORE the change is accepted as commit-ready. The discipline expands §11.4.4 (test-interrupt-on-discovery) + §11.4.8 (deep-web-research) + §11.4.43 (TDD-fix) + §11.4.50 (deterministic consistency) + §11.4.82 (Phase 1 forensic before any speculative patch) into an explicit pre-commit evaluation pipeline.

The 5-pass evaluation contract (mandatory, ALL passes must hold):

Pass 1 — Main-task verification. The change achieves the stated main task. Captured-evidence per §11.4.5 / §11.4.69 demonstrates the user-visible behaviour now works AS INTENDED. NO inferred / NO assumed / NO "should work" — only captured evidence.

Pass 2 — Regression-blast-radius analysis. Every file the change touches AND every file that imports / sources / references the changed file is enumerated. For each, the agent demonstrates (via grep / dependency-graph / actual test run) that the change does not break the existing contract. Examples: function signature changes → audit every caller; HAL config change → audit every test consuming the config; gate addition → audit every paired meta-test mutation. Captured evidence per §11.4.5.

Pass 3 — Cross-feature interaction analysis. Beyond direct-dependency blast radius, the change is evaluated against parallel features that share state, timing, hardware, or shell environment. Examples: audio HAL change → check video routing (Fix #88 SurfaceView Z-order timing); WiFi roaming change → check BT A2DP coex (Fix #70); test-script addition → check §11.4.84 working-tree quiescence + §11.4.89 background-test discipline. Captured evidence per §11.4.5.

Pass 4 — Deep-research validation. Per §11.4.8, the chosen approach is verified against external state-of-the-art (official docs, vendor guides, open-source codebases, kernel mailing list, AOSP gerrit, GitHub issues). Either the change cites an external precedent OR the literal "NO external solution found — original work" is recorded. CodeGraph queries per §11.4.78 + §11.4.79 surface intra-codebase precedents.

Pass 5 — Anti-bluff confirmation. Per §11.4 / §11.4.1 / §11.4.6 / §11.4.27 / §11.4.50 / §11.4.52 / §11.4.69 / §11.4.83 the change is verified to NOT introduce any new bluff surface — no metadata-only PASS, no config-only PASS, no script-bug FAIL-bluff, no "tests pass but feature doesn't work" anti-pattern. Captured evidence directories per §11.4.5 reference the exact passing paths.

Documentation requirement. Each Pass produces written documentation (commit-message footers OR `docs/` entries OR `qa-results/` evidence) demonstrating the pass was completed. A change without all 5 passes documented is severity-equivalent to a §11.4 PASS-bluff at the development-process layer.

Acceptance criteria. Only AFTER all 5 passes complete with documented evidence may the change be: (a) committed via `commit_all.sh`, (b) pushed via §11.4.88 background-push, (c) tested via §11.4.89 background test, (d) released / tagged per §11.4.40.

No escape hatch. No `--skip-multi-pass`, `--single-pass`, `--bluff-permitted`, `--fast-commit` flag exists. The 5 passes apply to every non-trivial change per the operator's verbatim mandate. Trivial changes (typo fixes, doc revision-header bumps, MD-export regeneration) are exempt ONLY when (i) the change touches zero source code AND (ii) the commit message explicitly cites the exemption.

Composes with §11.4 / §11.4.1 (anti-bluff baseline), §11.4.4 (test-interrupt-on-discovery — Pass 2 + Pass 3 trigger this), §11.4.5 (captured-evidence quality — every Pass's evidence MUST be quality-checked), §11.4.6 (no-guessing — Pass 1-3 require captured evidence not hypothesis), §11.4.8 (deep-research — Pass 4 expansion), §11.4.20 / §11.4.70 (subagent-driven — Pass 2-3 dispatchable to subagents per parallel analysis), §11.4.27 (no-fakes-beyond-unit — Pass 5 enforces), §11.4.42 (iteration discipline — multi-pass IS the iteration), §11.4.43 (TDD-fix — RED → Pass-1, GREEN → Pass-2-5), §11.4.50 (deterministic consistency — Pass 1 captured evidence MUST be N=3 reproducible), §11.4.52 (autonomous validation — Pass 5 anti-bluff confirmation gates on autonomous evidence), §11.4.69 (universal sink-side evidence — Pass 1+5 use closed-set taxonomy), §11.4.78 / §11.4.79 (CodeGraph — Pass 3 dependency-graph queries), §11.4.82 (Phase 1 forensic — §11.4.92 generalises Phase 1 to Pass-1+Pass-2+Pass-3 across every change class), §11.4.83 (docs/qa transcript — Pass evidence dir is the QA transcript), §11.4.85 (stress + chaos — Pass 2+3 evidence MUST include

stress + chaos signals when applicable), §11.4.87 (endless-loop — multi-pass runs concurrently across PWUs per §11.4.58), §11.4.89 (background tests — Pass 1-3 long tests run backgrounded).

Pre-build gate `CM-COVENANT-114-92-PROPAGATION` enforces this anchor literal across the canonical fleet. Pre-build gate `CM-MULTI-PASS-EVALUATION-EVIDENCE` audits recent commits for the 5-pass evidence trail (commit message footer OR per-commit `qa-results/<commit-hash>/passes/` directory OR `Issues.md / Fixed.md` narrative). Paired §1.1 meta-test mutations strip the load-bearing literals → gates FAIL.

Canonical authority: this `Constitution.md` §11.4.92 in the `HelixConstitution` submodule.

Non-compliance is a release blocker. A commit landing without the 5-pass evaluation documented is severity-equivalent to a §11.4 PASS-bluff at the development-process layer — it implies the change was vetted when it was not.

2026 05 27

§ . . — SQLite-backed single-source-of-truth for workable items (User mandate, - -)

Forensic anchor — verbatim user mandate (2026-05-27):

"There MUST be single source of truth for all of our workable items - SQLite database containing all workable items, all data about it and from there proper scripts (we recommend Go programs) which will update / regenerate all documentation - Issues, Fixed, summary docs, and everything related! We MUST reduce a chance for sync to be broken between the Db and all documents we have! We MUST BE able to generate always all docs from DB or to re-generate Db from all docs we have in opposite direction if we must! All this MUST BE covered with all supported test types, validated and verified by tests which will produce real proofs and be 100% anti-bluff compatible!"

The current text-based Issues.md / Fixed.md / Issues_Summary.md / Fixed_Summary.md / CONTINUATION.md / Status.md tracker constellation is converted to a **SQLite-database-backed single source of truth** with bidirectional regeneration between DB and Markdown. The DB is the authoritative source; all Markdown / HTML / PDF / Status / Summary surfaces are generator output. Sync drift is mechanically impossible because every regeneration starts from the DB.

Canonical artefacts (all consuming projects **MUST** adopt):

1. **Database file** at canonical path `docs/workable_items.db` (SQLite 3, **TRACKED in git, NEVER gitignored** per User mandate 2026-05-27 amendment). The DB IS the single source of truth — gitignoring it would defeat the SSoT mandate at the version-control layer. Every commit that mutates workable items **MUST** stage + push the DB file alongside the Markdown regen output. The DB is an exception to §11.4.30 build-artefact gitignore — it is NOT a build artefact, it IS authoritative source data. SQLite's WAL mode produces a sidecar `*.db-wal` + `*.db-shm` which ARE gitignored (transient WAL files; their state is checkpointed into the main `.db` file by `workable-items sync` or any explicit `PRAGMA wal_checkpoint(TRUNCATE)`).

2. Schema (mandatory minimum tables):

- `items` — primary key `atm_id` per §11.4.54; columns `type` (§11.4.16 closed-set), `status` (§11.4.15 + §11.4.90 closed-set including Obsolete), `severity`, `title`, `description` (≥ 6 words / ≥ 40 chars per §11.4.91), `created_at`, `last_modified`, `current_location` (Issues|Fixed), `forensic_anchor`, `closure_criteria`, `composes_with` (JSON array of refs).
- `item_history` — append-only audit log: `atm_id`, `event_type` (Opened|Updated|Reopened|Fixed|Implemented|Completed|Obsolete), `by` (AI|User per §11.4.34), `on_date`, `reason` (closed vocabulary per §11.4.34 + §11.4.90), `evidence_path`.
- `obsolete_details` — `atm_id`, `since`, `reason` (§11.4.90 closed-set), `superseding_item`, `triple_check_evidence` (mandatory per §11.4.90).
- `operator_block_details` — `atm_id`, `what`, `why_exhausted_alternatives`, `unblock_condition`, `who` per §11.4.21.
- `firebase_metadata` — per §11.4.47.

- `meta` — schema version, `last_sync_direction`, `last_sync_timestamp`, `integrity_hash`.

3. **Go binary** at canonical path `cmd/workable-items/` (separate Go module, NOT part of AOSP build). Subcommands:

- `workable-items sync md-to-db` — parse `Issues.md` + `Fixed.md` (single source of authority during transition window), upsert into DB. Idempotent.
- `workable-items sync db-to-md` — regenerate `Issues.md` + `Fixed.md` + `Issues_Summary.md` + `Fixed_Summary.md` + `CONTINUATION.md` §3 + every `Status.md` from DB. Idempotent.
- `workable-items diff` — print DB vs MD divergence. Used by pre-build gate.
- `workable-items validate` — schema sanity + §11.4.15/§11.4.16/§11.4.33/§11.4.34/§11.4.54/§11.4.90/§11.4.91 invariants per row. FAIL on any violation.
- `workable-items add <type> <severity> --title <title> --description <description>` — interactive entry with mandatory fields enforced.
- `workable-items close <atm-id> --status <fixed|implemented|completed|obsolete> --evidence <path>` — terminal transition with mandatory captured-evidence per §11.4.5/§11.4.90.

4. **Operational integration:**

- `commit_all.sh` pre-commit hook runs `workable-items diff` and refuses if MD/DB diverge.
- `sync_issues_docs.sh` invokes `workable-items sync db-to-md` instead of legacy bash generators (`generate_issues_summary.sh` becomes a shim wrapping the Go binary).
- `pre_build_verification.sh` runs `workable-items validate` as a gate; failure increments ERRORS.

5. **Anti-bluff test coverage** (mandatory per §11.4 + §11.4.85 + §11.4.27):

- Unit tests for every CRUD operation + every schema invariant.
- Integration test: full round-trip MD → DB → MD with byte-identical re-emission (modulo cosmetic-whitespace normalisation per closed-set tolerance).

- Stress + chaos per §11.4.85: 1000-row insert / concurrent-write (10 writers) / mid-write SIGKILL / corrupt-DB recovery / disk-full handling.
- Paired §1.1 meta-test mutation: strip a CRUD function → unit test FAILs.
- HelixQA Challenge entry `CME-WORKABLE-ITEMS-001` exercising end-to-end DB → MD → DB → MD round-trip.

Bidirectional regeneration guarantee. The DB MUST be reconstructible from the MD docs (for disaster recovery + audit replay) AND the MD docs MUST be reconstructible from the DB (the primary direction during normal operations). Either direction's regeneration produces semantically equivalent output (closed-set tolerance for cosmetic whitespace + section-ordering noise).

Cross-project propagation. Every consuming project that adopts this constitution submodule inherits the SQLite-SSoT approach. The Go binary lives in the constitution submodule (`constitution/scripts/workable-items/`) so consumers reference it from there per §11.4.74 catalogue-first discipline — never reimplement.

Migration path (per §11.4.42 iteration discipline + §11.4.58 PWU pipeline):

- **Phase 1:** file `§LA` Issues entry tracking the migration, scope-locked.
- **Phase 2:** Go binary scaffold + schema DDL committed.
- **Phase 3:** `sync md-to-db` lands + initial migration captures current state.
- **Phase 4:** `sync db-to-md` lands + byte-identical round-trip CI gate.
- **Phase 5:** existing generators (`generate_issues_summary.sh` etc.) become Go-binary shims.
- **Phase 6:** legacy text-direct edits prohibited (pre-commit hook).

Pre-build gates `CM-COVENANT-114-93-PROPAGATION` (anchor literal in canonical fleet) + `CM-WORKABLE-ITEMS-DB-PRESENT` (DB regen mechanism per §11.4.77 + DB schema-version current) + `CM-WORKABLE-ITEMS-MD-DB-IN-SYNC` (diff returns empty). Paired §1.1 meta-test mutations strip the load-bearing literals → gates FAIL.

Composes with §11.4 (anti-bluff covenant — DB is the captured-evidence registry), §11.4.12 (Issues_Summary sync — DB-driven), §11.4.15 (Status closed-set — DB schema enforces), §11.4.16 (Type closed-set — DB schema enforces), §11.4.17 (universal-vs-project — DB schema is universal), §11.4.19 (Fixed-document column-alignment — DB query emits both Summaries from same query),

§11.4.21 (Operator-blocked details — `operator_block_details` table), §11.4.27 (no-fakes-beyond-unit — DB integration tests exercise real SQLite), §11.4.30 (.gitignore + regeneration mechanism — DB file gitignored with `workable-items` sync as regen), §11.4.33 (closure vocabulary — DB enforces), §11.4.34 (Reopened source attribution — DB `item_history` table), §11.4.42 (iteration discipline — 6-phase migration), §11.4.43 (TDD — RED before each phase), §11.4.44 (revision header — Markdown output preserves), §11.4.45 (Status.md per integration — generator queries DB by domain), §11.4.50 (deterministic consistency — round-trip MD → DB → MD byte-identical), §11.4.51 (LIVE_ADB_FIRST — N/A, host-only), §11.4.52 (autonomous validation — DB integrity is autonomously verifiable), §11.4.53 (Fixed_Summary parity — DB ensures), §11.4.54 (ATM-NNN — DB primary key), §11.4.55 (Reopens-history doc — DB `item_history` queryable), §11.4.56 (Status_Summary parity — DB queries produce both pages), §11.4.57 (README doc-link section — DB drives), §11.4.58 (parallel PWU — Phase 1-6 land as separate PWUs), §11.4.60 (composite always-sync — DB → MD → HTML+PDF), §11.4.65 (universal MD export — final stage), §11.4.74 (catalogue-first — Go binary in constitution), §11.4.83 (docs/qa transcript — DB queries serve as audit replay), §11.4.85 (stress + chaos — DB integration test scope), §11.4.86 (roster/corpus auto-sync — DB schema accommodates roster tables), §11.4.87 (endless-loop — DB queries are zero-idle-friendly), §11.4.89 (background tests — DB validation runs detached), §11.4.90 (Obsolete status — `obsolete_details` table), §11.4.91 (Summary clarity — DB enforces description-floor at insert time), §11.4.92 (multi-pass evaluation — every DB schema migration follows 5-pass).

Canonical authority: this Constitution.md §11.4.93 in the HelixConstitution submodule.

Non-compliance is a release blocker. A consuming project that maintains text-based trackers without the DB-SSoT migration plan filed + Phase progression actively in flight is severity-equivalent to a §11.4 PASS-bluff at the data-architecture layer — the sync-drift surface §11.4.93 closes IS a §11.4 violation vector.

§ . . – Zero-idle priority-first parallel-by-default operating mode (User mandate, – –)

Forensic anchor — verbatim user mandate (2026-05-27):

"We MUST NEVER sit iddle / wait or sleep if there is possibility for us to work on something like in this case! Continue all work further in parallel if possible fully autonomously! We MUST not waste time since work scope we have is huge! ... Always check if there is a possibility to work on something while we are not working actively on something! Pick always by priority - most critical workable items and other tasks MUST BE done first! ... Everything we can should be done fully in background, should not affect stability of the System or create problems, should be done using subagents-driven approach! ... Stay still / iddle if nothing is left to be done at all or waiting for something that is blocking us / you!!!"

§11.4.94 is the **operating-mode reaffirmation** that binds §11.4.87 (endless-loop), §11.4.20 / §11.4.70 (subagent-driven), §11.4.42 (iteration-priority), §11.4.58 (parallel PWU), §11.4.72 (audio top-priority), §11.4.82 (iteration-speedup), §11.4.88 (background-push), §11.4.89 (background-test) into a SINGLE always-on enforcement contract:

The mandate (binding on the conductor + every subagent + every helper script):

(A) Idle-only-when-genuinely-blocked. The conductor MAY sleep / wait / be idle ONLY when (i) every priority-queued workable item is genuinely blocked on external dependency (operator hardware action, network upstream, hardware-attended forensic, build/test completion that the conductor cannot accelerate), OR (ii) the operator has issued an explicit STOP, OR (iii) host-session-safety demands it per §12. "I don't see what to do next" is NEVER a valid idle reason — the conductor MUST recheck the priority queue per §11.4.42 step 1 and §11.4.87(A) termination criteria.

(B) Always check parallel-work feasibility at every pause point. Before any wake-up schedule / before any "waiting for X" / before any sleep call, the conductor MUST: (1) survey the priority queue per §11.4.42 SCOPE LOCK + §11.4.72 audio-

first + §11.4.87(B) zero-idle, (2) identify ALL items that can progress without contending with the in-flight blocker, (3) dispatch them in parallel per §11.4.20 / §11.4.70 (subagent-driven if non-trivial) + §11.4.58 (PWU disjoint scope) + §11.4.89 (background long tests). Only AFTER step (3) returns "no parallel work available" may the conductor schedule a wake/sleep.

(C) Priority order MANDATORY at every pick. Per §11.4.42 priority axis + §11.4.72 audio-top-priority axis. Pick the highest-Severity / highest-priority item the conductor can autonomously progress. Lower-priority items wait until higher-priority queues drain or block on external dependency.

(D) Subagent-driven by default for non-trivial scope. Per §11.4.20 + §11.4.70. The conductor remains the integration + commit + push seam (per §11.4.58 + §11.4.84 quiescence + §11.4.88); subagents execute the analyses + scaffolds + tests + migrations in parallel. Inline execution permitted only when the task is trivial (single-file edit, sub-300-line diff) OR the operator explicitly requested it.

(E) Background by default for long-running work. Per §11.4.85 stress + §11.4.89 background test + §11.4.88 background push. Any operation expected to exceed ~30 s wall-clock MUST run detached (nohup + disown) with log under `qa-results/<op-id>_<ts>.log`. The conductor returns to the priority queue immediately and polls back when needed.

(F) Stability-preserving — no compromise. Parallel work MUST NOT compromise system stability or create new problems. Composes with §11.4.92 multi-pass change-evaluation — every parallel branch's output passes 5-pass before integration. Composes with §11.4.84 working-tree quiescence — concurrent subagents in same checkout coordinate via lockfile or git worktree. Composes with §12.6 / §12.7 / §12.8 / §12.9 host-session safety — parallel work respects 60% memory ceiling and -j2 AOSP cap.

(G) Status updates as work progresses. Per operator mandate "Let us know how goes this mandatory approach". The conductor surfaces parallel-work catalogue + in-flight branches + recent closures in compact summary form whenever the operator asks "how are we progressing" OR at natural milestone boundaries (after each constitution-anchor landing, after each commit batch, after each subagent return).

Anti-pattern explicitly forbidden:

- Scheduling a wake-up without first surveying the parallel-work queue per (B).
- Returning "nothing to do — sleeping" when non-trivial priority items remain that the conductor could autonomously progress.
- Serialising work that could run in parallel without justification.
- Picking lower-priority items while higher-priority ones remain progressable.
- Doing work foreground that should run via subagent or background.
- Letting parallel work introduce instability or regression.

Pre-build gate `CM-COVENANT-114-94-PROPAGATION` enforces this anchor literal across the canonical fleet. **Pre-build gate** `CM-PARALLEL-WORK-AUDIT` (when implemented) audits the last N wake-up schedule decisions for surfaced parallel-work-queue evidence — wake-ups without the survey are §11.4.94 violations. Paired §1.1 meta-test mutations strip the load-bearing literals → gates FAIL.

Composes with §11.4.20 (subagent-driven) / §11.4.42 (iteration discipline) / §11.4.58 (parallel-development PWU) / §11.4.70 (subagent-driven default) / §11.4.72 (audio top-priority) / §11.4.82 (iteration-speedup) / §11.4.84 (working-tree quiescence) / §11.4.85 (stress + chaos) / §11.4.87 (endless-loop zero-idle) / §11.4.88 (background-push) / §11.4.89 (background-test) / §11.4.92 (multi-pass evaluation) / §12.6 + §12.7 + §12.8 + §12.9 (host-session safety).

Canonical authority: this Constitution.md §11.4.94 in the HelixConstitution submodule.

Non-compliance is a release blocker. Conductor idle without exhausting the parallel-work queue is severity-equivalent to a §11.4 PASS-bluff at the project-execution layer — it implies the conductor has progressed the project when it has not. No escape hatch beyond the (A) genuinely-blocked / explicit-STOP / host-safety triad.

2026 05 27

§ . . — Workable-items SQLite DB is
TRACKED in git, NEVER gitignored (User mandate,
— —)

Forensic anchor — verbatim user mandate (2026-05-27):

"We shall not Git ignore our workable items SQLite DB since it is our single source of truth for current working items we have and all data related to it! We MUST fix this if DB is Git ignored! ... workable items SQLite DB regularly committed and pushed to all upstreams!"

The §11.4.93 workable-items SQLite database at `docs/workable_items.db` is the **AUTHORITATIVE source of truth** for every workable item in every consuming project. It MUST be:

- **TRACKED in git** at canonical path `docs/workable_items.db` . NEVER gitignored regardless of file-size or "build-artefact-class" heuristics — the DB is NOT a build artefact, it IS authoritative source data.
- **Committed alongside every workable-item state change.** Every `workable-items sync md-to-db` invocation that mutates DB state MUST stage + commit + push the DB file in the same commit as the corresponding MD changes per §11.4.19 atomic-move discipline.
- **Pushed to every upstream per §2.1** — every mirror of every consuming project carries the latest DB so a fresh clone from any mirror reconstructs the full workable-items inventory without consulting any other source.
- **WAL-checkpointed before commit** — SQLite's WAL mode produces transient `*.db-wal` + `*.db-shm` sidecars. Those ARE gitignored (per §11.4.30 transient/cache classification). The `workable-items` Go binary MUST execute `PRAGMA wal_checkpoint(TRUNCATE)` before any commit-stage step so the main `.db` file carries the authoritative state and the sidecars are safely discardable.
- **NEVER force-rewritten** without §9.2 hardlinked-backup + operator authorization. The DB's commit history IS the workable-items audit trail; rewriting it is data-loss equivalent to force-pushing a tracked Issues.md history.

§11.4.30 carve-out. This anchor is an explicit named exception to §11.4.30's build-artefact / data-file gitignore rule. The DB file at `docs/workable_items.db` is **NOT** a build artefact and is **NOT** transient data — it is the project's SSoT registry whose value increases monotonically as work lands. The §11.4.77 regeneration mechanism applies to other gitignored bulk data (`.git-backup-*` , `RKTools/linux/` , etc.) — NOT to this DB.

§11.4.93 amendment. §11.4.93's earlier text describing the DB as "gitignored per §11.4.30 with §11.4.77 regeneration mechanism" is hereby **AMENDED**: the DB is TRACKED, not gitignored. The Markdown trackers (Issues.md / Fixed.md / Summaries / Status.md fleet / CONTINUATION.md) remain tracked AND derived; the DB is tracked AND authoritative. Both directions of regeneration (`workable-items sync db-to-md` + `workable-items sync md-to-db`) continue to land per §11.4.93's bidirectional round-trip guarantee.

Pre-build gates `CM-COVENANT-114-95-PROPAGATION` (anchor literal across canonical fleet) + `CM-WORKABLE-ITEMS-DB-TRACKED` (`git ls-files docs/workable_items.db` returns non-empty if DB exists on disk + `.gitignore` does NOT exclude it). Paired §1.1 meta-test mutation adds the DB to `.gitignore` → gate FAILs.

Composes with §11.4.93 (SQLite-SSoT — §11.4.95 amends its tracking-policy clause), §11.4.30 (`.gitignore` discipline — explicit carve-out anchored here), §11.4.77 (regeneration mechanism — does NOT apply to the DB; the DB IS the source), §2.1 (multi-upstream push — DB pushed to every mirror), §9.2 (data safety — destructive DB ops require hardlinked-backup + operator authorization), §11.4.41 (force-push merge-first — applies if DB history ever needs rewrite).

Canonical authority: this Constitution.md §11.4.95 in the HelixConstitution submodule.

Non-compliance is a release blocker. A consuming project that gitignores its workable-items DB is severity-equivalent to a §11.4 PASS-bluff at the source-of-truth layer — it implies SSoT exists when it is in fact ephemeral local-only state.

2026 05 27

§ . . — Safe-parallel-work-with-long-build catalogue + mandate (User mandate, — —)

Forensic anchor — verbatim user mandate (2026-05-27):

"Are there except AOSP build process any other active jobs being done at the moment? Can we work on something in parallel while build is in progress so we slowly cleanup our slate? Anything workable which will not affect the build process or final image we generate for the flashin. If yes, and that was not already clear by our root Constitution, we should add all mandatory details into it ... Make sure we do as much as possible work in background in parallel with main work stream and oreferably using subagents-driven approach!"

§11.4.94 mandates "always check parallel-work feasibility before idle"; §11.4.89 mandates background-test discipline; §11.4.58 defines the parallel-work-unit pipeline. §11.4.96 is the **operational catalogue** that enumerates, for the canonical long-running workload (5-7 h AOSP containerised build per §12.9), exactly what CAN safely run in parallel and what MUST NOT — closing the ambiguity gap §11.4.94 left implicit.

Closed-set categories — SAFE during AOSP containerised build:

- **(A) Markdown / documentation work** anywhere under `docs/` , `*.md` , `README`, `CLAUDE.md`, `AGENTS.md`, `QWEN.md`, `Status.md` / `Status_Summary.md`, `Issues.md` / `Fixed.md` / their Summary siblings, changelogs, research notes. The build container mounts source read-only at a snapshot moment + reads only the build-relevant files (`Android.bp` / `Android.mk` / source code / vendor binaries). MD edits do NOT propagate into the build's read scope.
- **(B) Generator + helper script work** under `scripts/` , `scripts/testing/` , `scripts/llm/` , `scripts/firebase/` , etc. — these are host-side helpers NEVER consumed by the AOSP build. Safe to edit + commit + push freely.
- **(C) Pre-build / meta-test gate authoring + paired §1.1 mutations** under `device/rockchip/rk3588/tests/pre_build_verification.sh` + `scripts/testing/meta_test_false_positive_proof.sh` — these run host-side via `bash pre_build_verification.sh` ; the AOSP build itself does NOT execute them inline. Safe to edit during build; the NEXT pre-build run (post-current-build) picks up the new gates.

- **(D) On-device tests** under `device/rockchip/rk3588/tests/test_*.sh` — these run on D3/D4 via ADB post-flash. Authoring new test scripts during build does NOT affect the build (the build packages the test corpus from the source tree but does so once at copy-out time; scripts added after the copy-out are absent from THIS image but present in source for the NEXT image).
- **(E) Constitution submodule edits + push** — separate git workspace at `constitution/`, not on the AOSP build's path. Edits + push to all 6 constitution remotes safe at any time.
- **(F) Project submodule commit + push** to either the parent or owned submodules' remotes — per §11.4.88 background push, this is OS-level orthogonal to the container.
- **(G) D3 / D4 live-ADB probes** (read-only `dumpsys` / `getprop` / `cat /proc/asound/...` / `screencap` / `logcat -d`) — these query device state without changing it; do NOT affect the host build.
- **(H) Subagent dispatch** per §11.4.20 / §11.4.70 — each subagent runs in isolated context window; multiple subagents in same checkout coordinate via §11.4.84 working-tree quiescence (lockfile or git worktree per disjoint scope). Read-only analysis subagents (obsolescence audit, summary clarity sweep, code-review) are uniformly safe.
- **(I) Web research + external API queries** (with §11.4.10 credential discipline preserved) — operator-relevant browsing, ASUS QVL lookups, vendor spec verification, package-registry version checks.
- **(J) Workable-items DB operations** per §11.4.93 + §11.4.95 — `workable-items sync` (when implemented), `workable-items validate`, `workable-items report` — read/write SQLite at `docs/workable_items.db`, orthogonal to build.
- **(K) Pre-build verification + meta-test execution** — RUN per §11.4.89 background-test discipline. The build container reads source; pre-build runs host-side reading the same source. No conflict at filesystem layer (read-only intersection).

Closed-set — UNSAFE during AOSP containerised build:

- **(α) git checkout / git reset --hard / git clean -df** on the source tree — would change files the build is reading; introduces inconsistency. Defer until build completes OR use a `git worktree` per §11.4.84.

- **(β) Mass file deletions or renames** under `device/` , `frameworks/` , `hardware/` , `vendor/` , `kernel-5.10/` , `external/` , `packages/` , `bionic/` , `bootable/` , `build/` , `system/` , `cts/` , `tools/` , `prebuilts/` , etc. — anything the AOSP build potentially reads. Apply ONLY after build completes.
- **(γ) Submodule pointer updates that affect built APKs** — bumping `device/rockchip/atmosphere/presenter` or `smarttube-player` or `vlc-player` etc. pointer changes the binary content the build packages. Defer pointer bumps to between-build windows OR plan for the NEXT build to pick them up.
- **(δ) out/ directory mutations** — the build's output directory. Never touch concurrently.
- **(ε) make clean / m clobber / rm -rf out/** while build runs — direct conflict; will likely crash the build.
- **(ζ) Container destruction** (`podman stop atmosphere-aosp-build` , `podman pod rm`) — terminates the build mid-flight. Operator-explicit only.
- **(η) Disk-filling operations** (large file downloads, container image builds, NVMe duplication) that push the host below the §12.9 free-space minimum — could OOM the build's writable layer.
- **(θ) Host-session-safety breaches** per §12 — suspend / hibernate / logout / unbounded memory ops in user.slice — kill the build's parent cgroup.

Conductor responsibility (per §11.4.94 already; §11.4.96 makes the catalogue explicit):

- Before EVERY pause point during a long build, the conductor MUST consult this catalogue, identify all (A)-(K) items in the priority queue per §11.4.42 + §11.4.72, dispatch at least one per §11.4.20 / §11.4.70 subagent-driven default + §11.4.89 background-task discipline. "Build is running, nothing else to do" is NEVER true per §11.4.94 + this catalogue.

Subagent-driven default for catalogue items:

- (A) MD work + (E) constitution edits + (F) commits: conductor-direct (lightweight + critical-state sequencing).
- (B) generator/helper work + (C) gate authoring + (D) on-device test authoring + (G) live-ADB probes + (I) web research + (J) DB ops: subagent-dispatchable per §11.4.20.

- (H) is itself the subagent dispatch — composes with the rest.
- (K) pre-build / meta-test execution: backgrounded per §11.4.89, conductor polls outcome.

Pre-build gate `CM-COVENANT-114-96-PROPAGATION` enforces this anchor literal across the canonical fleet. Pre-build gate `CM-PARALLEL-WORK-DURING-BUILD-AUDIT` (when implemented) audits recent commits during AOSP build windows for the parallel-work output (commits landing during a build window with timestamps inside `qa-results/aosp_build_*.log` start/end). Paired §1.1 meta-test mutations strip the catalogue literals → gates FAIL.

Composes with §11.4.20 (subagent-driven) + §11.4.42 (iteration priority) + §11.4.58 (parallel PWU) + §11.4.70 (subagent default) + §11.4.72 (audio top-priority) + §11.4.82 (iteration-speedup) + §11.4.84 (working-tree quiescence — esp. for subagent parallelism) + §11.4.85 (stress + chaos can run in parallel) + §11.4.87 (endless-loop zero-idle) + §11.4.88 (background-push) + §11.4.89 (background-test) + §11.4.92 (multi-pass evaluation per parallel branch) + §11.4.93 / §11.4.95 (workable-items DB) + §11.4.94 (parallel-work survey mandate) + §12.6 / §12.7 / §12.8 / §12.9 (host-session safety + build cap).

Canonical authority: this Constitution.md §11.4.96 in the HelixConstitution submodule.

Non-compliance is a release blocker. Conductor scheduling a wake/sleep during an active AOSP build without first dispatching at least one (A)-(K) catalogue item from the priority queue is a §11.4.94 + §11.4.96 violation = §11.4 PASS-bluff at the operating-mode layer.

2026-05-27

**§ . . — Maximum-use-of-idle-time
mandate + progress-update cadence (User mandate,
- -)**

Forensic anchor — verbatim user mandate (2026-05-27):

"keep it working, we should do as much as possible, if not it all but as much as we can as long as there is idle time! it MUST be used! make sure that it is clear in constitution! keep us updated about all progress and all physical proofs and gathered data as you progress through all open workable items!"

§11.4.97 strengthens §11.4.87 (endless-loop) + §11.4.94 (zero-idle priority-first) + §11.4.96 (safe-parallel-with-build catalogue) with an **explicit maximum-use-of-idle-time mandate** AND a **progress-update cadence mandate**:

(A) Maximum-use idle-time mandate. Every minute of conductor idle time during which (i) the conductor has work it could autonomously progress AND (ii) is NOT genuinely blocked on external dependency (operator hardware action, network upstream, host-session safety, build/test completion) is a §11.4.97 violation. "As much as possible, if not it all but as much as we can" is the operative phrasing — the conductor MUST dispatch work continuously through the entire idle window, not just sample-check at scheduled wakes.

(B) Progress-update cadence mandate. The conductor MUST emit operator-facing progress updates at the following natural milestone boundaries (no operator prompt required):

- Every commit that lands (HEAD advance) — 1-line "what just landed".
- Every subagent return — 1-line "subagent X returned, integrated".
- Every constitutional anchor landed — 1-line "§X.Y.Z propagated to 6 remotes".
- Every gathered physical proof — 1-line "captured evidence at qa-results/".
- Every milestone closure (Issues → Fixed migration, Obsolete classification, etc.) — 1-line.

Progress updates are CONCISE (1-3 lines each); the operator's mental model stays current without conductor overhead.

(C) Continuous physical-proof gathering mandate. Per §11.4.5 + §11.4.6 + §11.4.69 — every autonomous-fixable item the conductor closes MUST gather positive captured-evidence: ADB probes, dumphsys captures, screencaps, ALSA hw_params, sysfs reads, file-content hashes, runtime metric snapshots. The

evidence is committed alongside the closure narrative. Per §11.4.93 SQLite-SSoT, the evidence path goes into the `item_history.evidence_path` column when DB integration lands.

(D) Per-anchor composition. §11.4.97 binds together: §11.4.5 (captured-evidence quality), §11.4.6 (no-guessing — every progress claim cites evidence), §11.4.13 (sink-side evidence when applicable), §11.4.20 (subagent-driven default for parallel scope), §11.4.27 (no-fakes-beyond-unit — closures cited with real captured evidence), §11.4.42 (priority queue), §11.4.50 (deterministic consistency — N=3 where applicable), §11.4.52 (autonomous validation), §11.4.69 (sink-side taxonomy), §11.4.70 (subagent-driven), §11.4.72 (audio top-priority), §11.4.83 (docs/qa transcript), §11.4.85 (stress + chaos), §11.4.87 (endless-loop), §11.4.88 (background-push), §11.4.89 (background-test), §11.4.94 (zero-idle), §11.4.96 (safe-parallel catalogue).

(E) Idle-only-when-genuinely-blocked. Same closed-set as §11.4.94(A): operator STOP, external dependency the conductor cannot accelerate, host-session-safety demand. Per §11.4.96 the safe-during-build catalogue (A-K) covers what's progressable; the conductor MUST exhaust the catalogue before scheduling sleep.

Operator-facing visibility. §11.4.97's progress-update cadence is the contract by which the operator stays informed without prompting. The conductor does NOT wait for operator status checks — milestone updates emit autonomously.

Pre-build gate `CM-COVENANT-114-97-PROPAGATION` enforces this anchor literal across the canonical fleet. Pre-build gate `CM-IDLE-TIME-AUDIT` (when implemented) audits the conductor's wake/sleep schedule across recent sessions for unjustified idle periods. Paired §1.1 meta-test mutations strip the load-bearing literals → gates FAIL.

Composes with every anchor in §11.4 (this is the operating-mode capstone):
§11.4.5 + §11.4.6 + §11.4.13 + §11.4.20 + §11.4.27 + §11.4.42 + §11.4.50 +
§11.4.52 + §11.4.69 + §11.4.70 + §11.4.72 + §11.4.83 + §11.4.85 + §11.4.87 +
§11.4.88 + §11.4.89 + §11.4.94 + §11.4.96.

Canonical authority: this Constitution.md §11.4.97 in the HelixConstitution submodule.

Non-compliance is a release blocker. Conductor idle when progressable work remains, OR conductor failing to emit milestone progress updates, is severity-equivalent to a §11.4 PASS-bluff at the operating-mode layer.

2026 05 28

§ . . — Full-Automation Anti-Bluff
Mandate — Live tests **MUST** be re-runnable end-to-end without manual intervention (User mandate, — —)

Forensic anchor — verbatim user mandate (2026-05-28):

"Make sure we have full automation testing of all scenarios with real bot, main group and users without any manual intervention or contribution of real user! Everything **MUST BE** fully automatic and autonomous! These tests **MUST BE** able to rerun endless times when needed! This is important to be done like this! It is critical! Continue all work and make this happen! We need such full automation testing so the whole System **MUST BE** fully validated and verified before it is integrated to our main projects! Make sure there is no false positives in testing! Every test and its results **MUST** obtain real proofs of everything working! No bluff is allowed! **IMPORTANT:** Make sure that all existing tests and Challenges do work in anti-bluff manner - they **MUST** confirm that all tested codebase really works as expected!"

§11.4.98 composes with §11.4 / §11.4.2 / §11.4.5 / §11.4.50 / §11.4.85 / §11.4.87 / §11.4.89 / §11.4.94 — closes the **manual-intervention gap** that they did not explicitly forbid. A live/integration/e2e/Challenge test which requires a human action during execution (typing a chat message, clicking a UI, hand-triggering a webhook, manually attaching a file, anything beyond the test starting up and reporting PASS/FAIL on its own) is **by definition a §11.4 PASS-bluff at the automation layer**, regardless of how thorough the manual run is. The reason: such a test cannot run continuously in CI, cannot validate regressions between manual runs, and the human dependency masks any drift between code changes and what the test exercises.

(A) Binding rule. Every test that this Constitution governs — unit / integration / e2e / Challenge / stress / chaos / live — MUST be fully self-driving end-to-end. The test process must accept inputs from configuration alone, exercise the load-bearing code path autonomously, and report PASS/FAIL/SKIP-with-reason without any further human action.

(B) Single permissible exception — one-time credential bootstrap.

Configuration performed OUTSIDE test execution is acceptable: populating `.env` from a vault, exporting shell env vars in `~/.bashrc`, OAuth approval at first install, MTPROTO session activation at first run. This is configuration, not test driving. Once credentials are present, the test MUST run end-to-end without any further human action. The configuration step is itself a §11.4.10 credential-handling concern, not a §11.4.98 automation concern.

(C) Concrete requirements for live messenger / channel / agent tests (the surface that exposed this mandate):

1. **No "operator MUST type a message" prompts in test scripts.** Tests requiring inbound messages MUST drive those messages programmatically — via a separate user account (MTPROTO for Telegram, real-user-token API for Slack, IMAP-test-account for email, etc.), via a webhook fixture, or via an in-process loopback. Never via human keystrokes during test execution.
2. **No hard-coded session UUIDs that collide with the active dev session.** Test runs MUST use a dedicated test-only session that the production / dev session is NOT simultaneously using. (Lesson from Herald 2026-05-28: `claude --resume <UUID>` invoked against the same session ID the dev session is using returns exit -1 with no output — a silent collision that fails for non-obvious reasons.)
3. **No 60-second human-response windows.** These create flake/timing races that are themselves §11.4.50 determinism violations: a single test invocation's PASS/FAIL depends on a human reaction-time variable, not on the code under test.
4. **Re-runability proof.** Every live test MUST PASS at least 3 consecutive automated invocations (`-count=3` for Go, `repeat 3` equivalent for other runtimes) with no state cleanup between runs. Persistent side effects (chats, files, DB rows, attached objects, queued events) MUST be cleaned by the test itself in a `defer / teardown` block.

5. **§11.4.98 obsolescence audit.** Every existing test that ships under this Constitution MUST be classified as either "self-driving" (COMPLIANT) or "manual-dependency" (NON-COMPLIANT — must be rewritten or marked §11.4.90 Obsolete). The audit is a release-gate item — a release that ships a NON-COMPLIANT test still pretending to PASS is blocked.
6. **No false-positive PASS.** A test PASS that does NOT exercise the load-bearing code path — e.g. a "live" test that skips silently because env vars are missing yet reports PASS, or a "live" test where the captured evidence is from a prior cached run — is itself a §11.4 PASS-bluff at the gating layer. SKIP-with-reason is the §11.4.3 correct posture; SKIP-reported-as-PASS or stale-evidence-reported-as-fresh is forbidden.

(D) Composition. §11.4.85 (stress + chaos) + §11.4.89 (background-test execution) + §11.4.87 (endless-loop autonomous work) + §11.4.94 (zero-idle parallel-by-default) + §11.4.98 (full-automation) together make this Constitution's testing surface a continuously-validated, fully-automated, non-flake, anti-bluff regime. Each closes a different gap; remove any one and the whole property collapses. §11.4.98 specifically closes the manual-intervention gap: §11.4 + §11.4.85 + §11.4.89 + §11.4.87 + §11.4.94 do not forbid a test from requiring human action, only from skipping the load-bearing path. §11.4.98 makes the human dependency itself a §11.4 violation.

(E) Inheritance per §11.4.35. §11.4.98 propagates to every consuming repository via the existing inheritance gate: each project's CLAUDE.md / AGENTS.md / QWEN.md MUST carry a short-form restatement citing the literal anchor `11.4.98`. The pre-build gate `CM-COVENANT-114-98-PROPAGATION` (when implemented) enforces this literal anchor presence across the canonical fleet. Paired §1.1 meta-test mutations strip the load-bearing literal → gates FAIL.

(F) Enforcement. A commit that adds or modifies a test that requires manual human action during execution is blocked at release-gate. A test classified as "manual-dependency" in the §11.4.98 audit that has not been rewritten within 30 days of classification graduates to §11.4.90 Obsolete and is removed from the active test suite (not deleted — preserved with `Obsolete-Details:` per §11.4.90, citing §11.4.98 as the obsolescence reason).

Canonical authority: this Constitution.md §11.4.98 in the HelixConstitution submodule. All consuming projects (Herald, a consuming project, future) restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker. No `--manual-test-OK` , `--skip-114-98-audit` , `--bluff-tolerance-temporary` flag exists. The 2026-05-28 user mandate is unambiguous: "No bluff is allowed!"

§ . 2026 05 28 — Latest-Source Documentation Cross-Reference Mandate — instructions, guides, and manuals **MUST** be verified against the latest official online sources **BEFORE** publication (User mandate, — —)

Forensic anchor — verbatim user mandate (2026-05-28):

"Make sure we ALWAYS check against latest versions of services we use web / online docs before creating instructions! This situation is illustration of how we can misguide ourselves or get banned! Add this all important generic / general points as proper mandatory rules when we are creating documentation or guides! These are mandatory rules / constraints and the result is consistency and safety of created instructions, guides and manuals!"

Case study that anchored this rule (Herald 2026-05-28). The first-draft Herald MTPProto setup guide (committed at `35fc10c` / `fb00354` / `f089dd6`) recommended VoIP / Google Voice / Twilio / TextNow numbers as a budget-friendly fallback option AND omitted the critical `recover@telegram.org` pre-login email step. Both pieces of guidance directly contradicted (a) Telegram's official documentation at https://core.telegram.org/api/obtaining_api_id ("all accounts that sign up or log in using unofficial Telegram clients are automatically put under observation") AND (b) the gotd/td maintainer's "How to not get banned?" guide (vendored at `submodules/gotd-td/.github/SUPPORT.md`). An operator who followed the original guide could have had their Telegram account permanently banned with no appeal — the very harm the documentation was supposed to help them avoid. The corrected guide landed at `8470ba7` after a forced cross-reference

against the latest official sources. This case study is forensic evidence that misguidance-by-stale-docs is the same severity class as a §11.4 PASS-bluff at the documentation layer.

Composition. This anchor composes with §11.4.4 (test-interrupt-on-discovery), §11.4.5 (captured-evidence), §11.4.8 (deep-research validation), §11.4.92 (multi-pass change-evaluation Pass 4), §107 (anti-bluff), and §11.4.98 (full-automation). It closes a specific gap §11.4.92 Pass 4 alludes to but does not explicitly mandate: **every operator-facing instruction, guide, manual, troubleshooting cookbook, or setup walkthrough MUST be cross-referenced against the LATEST official online documentation of the service / library being documented BEFORE the document is committed.**

(A) Binding rule — pre-commit cross-reference. Before committing any document that contains operator-actionable steps (setup walkthroughs, integration guides, API how-tos, credential acquisition, security configurations, troubleshooting cookbooks, billing-policy documentation, terms-of-service summaries, any external-service interaction guide), the author MUST:

1. **Fetch the latest official online documentation** of the third-party service / library being documented via WebFetch, direct browsing, the service's own MCP server (if available), or equivalent authoritative real-time source. Do NOT rely on training data, memory, prior assumptions, or older committed docs as the source of truth.
2. **Cross-reference each instruction in the new document** against that source. For each step the operator will take, verify: (a) the service still supports that action; (b) the form fields / parameters / endpoints are still the same; (c) any new constraints, deprecations, or warnings the service published since the doc's last verification.
3. **Seek secondary authoritative sources when the official documentation is sparse / silent** on a critical requirement. Examples: the library maintainer's official guides (`submodules/<lib>/README.md` , `SUPPORT.md` , `SECURITY.md`), the service's official blog / changelog, the service's official support email / channel responses, community-vetted FAQs (Stack Overflow accepted answers, official Discord pinned messages).

4. **Cite the source URL + the date checked at the bottom of the document** in a `## Sources verified` section. Example: `Sources verified 2026-05-28: https://core.telegram.org/api/obtaining_api_id + submodules/gotd-td/.github/SUPPORT.md` . The citation is the audit trail; a doc without it is by §11.4.99 definition NOT verified.
5. **Cite the cross-reference in the commit message footer.** A commit that adds or modifies operator-facing instructions MUST include a `Sources verified <date>: <urls>` footer line so the verification trail is reachable from `git log` queries.

(B) Negative-finding documentation is required. If the cross-reference reveals the official source is silent / contradictory / outdated on a question the new document needs to answer, the document MUST explicitly note that gap so the next reader does not assume the absence of contradiction means authoritative agreement. Example phrasing the Herald MTPROTO guide uses: *"The official Telegram docs do not enumerate App title validation rules; the constraints below were verified against the gotd/td community channel + empirical operator testing 2026-05-28."*

(C) Re-verification cadence. Documents older than 6 months without re-verification are STALE — operators MUST NOT trust them for fresh action without re-running the cross-reference. Documents MUST be re-verified:

1. **Before being cited as the authority for an operator-action campaign** (the operator's command to "follow this guide" is the trigger).
2. **At every major release boundary** (vN.0.0 tag — the release-gate sweep includes a §11.4.99 freshness audit).
3. **Whenever the documented service publishes a breaking-change announcement** (the agent that knows about the announcement MUST queue the re-verification).
4. **When an operator reports an error following the guide** (the case-study anchoring is automatic and triggers immediate re-verification).

(D) Service-specific risk-classifications. Documentation for the following service families MUST include explicit safety warnings cross-referenced against the latest published policies, with a §11.4.99 "Sources verified" date NEVER older than 90 days:

Family	Specific risks the documentation MUST address
Telegram / WhatsApp / messenger APIs (unofficial clients)	Anti-abuse-system observation; one-phone-one-app-id limits; ban-on-VoIP policies; rate-limit floods; user-impersonation risks; pre-login declaration emails.
Cloud provider APIs (AWS, GCP, Azure, Cloudflare)	Billing blast-radius; account-lock policies; service-quota limits; least-privilege IAM; data-residency obligations.
Payment systems (Stripe, PayPal, banks, Mercado Pago)	KYC/AML compliance; PCI-DSS data-handling rules; webhook signature verification; idempotency keys; refund/chargeback flows.
AI / LLM providers (Anthropic, OpenAI, Google AI)	Terms-of-service violation triggers; rate-limit + quota policies; data-retention defaults; safety-classifier outcomes; PII-in-prompt risks.
Code-hosting services (GitHub, GitLab, GitFlic, GitVerse, Bitbucket)	Token-leak revocation policies; force-push to protected-branch policies; rate-limit-on-API-tokens; secret-scanning bot behaviour.
OS / package managers (apt, brew, npm, pip, cargo)	Supply-chain compromise vectors; signature verification; lockfile discipline; mirror-trust policies.

This list is **NOT exhaustive** — when a new external service is documented, the §11.4.99 author judges whether the service has comparable risk surface; if yes, the same safety-warning requirement applies.

(E) Composition with §11.4.92 multi-pass change-evaluation. §11.4.92 Pass 4 (deep-research validation) already requires "external precedent or literal 'NO external solution found'". §11.4.99 strengthens this specifically for the operator-facing-instructions class: deep research is mandatory for ALL such docs, NOT just non-trivial code changes. The agent that authors an instruction guide CANNOT cite §11.4.92 Pass 4 as a substitute for §11.4.99 — the two pass independently.

(F) Inheritance per §11.4.35. Every consuming repository's CLAUDE.md / AGENTS.md / QWEN.md MUST carry a short-form restatement citing the literal anchor 11.4.99 . The pre-build gate CM-COVENANT-114-99-PROPAGATION (when implemented) enforces this literal anchor presence across the canonical fleet. Paired §1.1 meta-test mutations strip the load-bearing literal → gates FAIL.

(G) Enforcement. A commit that adds or modifies operator-facing instruction documentation without (a) a "Sources verified " footer in the document itself AND (b) a Sources verified line in the commit-message footer is blocked at release-gate. A document with operator-actionable steps that becomes stale (older than 6 months without re-verification, or 90 days for risk-classified services per §(D)) graduates to §11.4.90 Obsolete with Obsolete-Details: Reason=stale-documentation; superseded-by=<replacement-doc-or-rewrite> after the 30-day grace window per §11.4.90.

Canonical authority: this Constitution.md §11.4.99 in the HelixConstitution submodule. All consuming projects (Herald, a consuming project, future) restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker. No --skip-source-check / --documentation-freshness-optional / --cite-sources-later / --trust-prior-doc-as-authoritative flag exists. The 2026-05-28 user mandate is unambiguous: the result MUST be "consistency and safety of created instructions, guides and manuals".

§11.4.100 — RETIRED. Demoted to consumer project (a consuming project's video-color/visual-quality fidelity) per §11.4.17/§11.4.35 — project-specific (RK3588/MPV/Arvus), not universal. See the consuming project's Constitution/CLAUDE/AGENTS/QWEN.

2026 05 28

§ . . — Autonomous-decision-over-blocking mandate (User mandate, - -)

Short tag: autonomous-decision-over-blocking .

Forensic anchor — verbatim user mandate (2026-05-28):

"when working in endless working loop fully autonomously try to decide most properly about points which would block execution and wait for us. If we haven't answered now work would be blocked whole night! If possible and if that will not cause any issues make proper and most reliable and safe decision so we achieve maximal efficiency and work gets fully done!"

The mandate. When operating in an autonomous / endless-loop mode (per §11.4.87), the agent **MUST** minimize operator-blocking and instead make the safe, reliable, reversible decision itself — so work is NOT stalled (e.g. overnight) waiting for input. The §11.4.87 endless-loop covenant told the agent to keep working; §11.4.101 tells it **HOW** to handle the decision points that would otherwise force it to stop and wait. A wrong bias toward blocking-when-it-was-safe-to-decide wastes operator time (the forensic harm: "blocked whole night"); a wrong bias toward deciding-when-it-should-have-blocked risks irreversible harm. The closed-set decision rule below is the precise boundary between the two.

Decision rule (closed-set — the agent MAY proceed autonomously when ALL hold):

(a) the action is **reversible** OR has a captured pre-op backup per §9.2 (hardlinked `.git` mirror, recorded refs/tags/submodule pointers); (b) the agent can **determine the safe choice from captured evidence** per §11.4.6 (no guessing — `LIKELY` / `probably` / `seems` is NOT a safe-choice determination; the evidence must state the correct choice as fact, or the condition is not met); (c) a wrong choice's **blast radius is bounded AND recoverable** — the damage is local, undoable, and does not cascade beyond the work unit; (d) it **composes with the existing covenants** — anti-bluff §11.4 (no PASS-bluff to skip the decision), host-session-safety §12 (no decision that suspends / logs-out / overruns the §12.6 memory budget), data-safety §9 (no destructive op without the §9 protocol).

Block-only-when rule (the agent MUST BLOCK — ask via the §11.4.66 interactive mechanism — ONLY when ALL of the following hold): the action is **irreversible AND high-blast-radius AND the safe choice cannot be determined from evidence**. Canonical block-required examples: external-account state the agent cannot inspect (registration, certification, billing portals), hardware the agent cannot physically access, destructive repository ops without a backup, force-push (which independently requires §9.2 + §11.4.41 authorisation), spending money or

sending messages / data to third parties. `operator-blocked` per §11.4.21 remains the last-resort status and is reached only after this rule fires AND the §11.4.21 self-resolution-exhaustion audit is complete.

Maximize-progress-while-blocked rule. When a block IS unavoidable, the agent MUST still maximize progress on every NON-blocked item in parallel per §11.4.87 + §11.4.94 (zero-idle priority-first) rather than idling — the blocked decision parks one work unit, it does not pause the loop. Posing the §11.4.66 question and continuing other work is the correct behaviour; posing the question and going idle is a §11.4.94 + §11.4.97 violation.

Composes with §11.4.6 (no-guessing — sub-rule (b) IS the no-guessing test applied to the decision itself), §11.4.21 (operator-blocked is the last-resort status the block-rule routes into, after self-resolution exhaustion), §11.4.40 (a release-tag decision is irreversible-and-high-blast-radius → always blocks per the block-rule), §11.4.41 (force-push is an explicit block-required example), §11.4.66 (the interactive-clarification mechanism the block-rule uses), §11.4.87 (the endless-loop covenant §11.4.101 refines — §11.4.87 says keep working, §11.4.101 says how to clear the decision points), §11.4.94 (zero-idle — maximize-progress-while-blocked), §9.2 (backup is the reversibility substitute in sub-rule (a)), §12 (host-safety bounds every autonomous decision).

Classification: universal (§11.4.17) — autonomous-decision-vs-block boundary is a reusable discipline for ANY project running an agent in an endless / unattended loop, hardware- and domain-agnostic.

4-layer coverage per §11.4.4(b). Propagation gate `CM-COVENANT-114-101-PROPAGATION` enforces the literal anchor `11.4.101` across the canonical consumer fleet (parent + owned-submodule `CLAUDE.md` / `AGENTS.md` / `QWEN.md`). Paired §1.1 meta-test mutation strips the `11.4.101` literal from a consumer file → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this `Constitution.md` §11.4.101 in the `HelixConstitution` submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--always-block-on-decision` , `--never-decide-autonomously` , `--skip-decision-rule` , `--block-without-self-resolution` flag exists. The

2026-05-28 user mandate is unambiguous: in autonomous mode the agent makes the safe reversible decision and keeps working — it does NOT block the operator overnight on a choice it had the evidence to make safely.

2026 05 29

§ . . — Mandatory systematic-debugging activation + always-loaded skill-discovery + plugin-dependency availability (User mandate, - -)

Short tag: `systematic-debugging-always-on` .

Forensic anchor — verbatim user mandate (2026-05-29):

"Make sure that we ALWAYS trigger / start the "/superpowers:systematic-debugging" skills when any issues happen! If this is possible to activate and use in this situations out of the box when we spot problems / issues / bugs / misalignments / inconsistencies we MUST activate the skill(s) and make strongest efforts in full in depth analysis / debugging and determine root causes of all problem or obtain relevant data and information we need! ... we MUST make sure that "/using-superpowers" skill is ALWAYS loaded, applied and used! All dependencies (plugins) that Claude Code or other market places are offering MUST BE installed if these are not already available for loading and use!"

This anchor binds three cooperating invariants. Together they make the difference between an agent that **guesses-and-retries** (symptom-patching, "transient", "flaky", "probably a race") and an agent that **investigates-to-root-cause-first** mechanical, every single time a problem surfaces. The §11.4.6 no-guessing mandate said *do not state an unproven cause*; §11.4.102 says *the moment a problem appears, the very first action is to launch the structured root-cause investigation skill — not to propose a fix*.

(A) Mandatory systematic-debugging activation on ANY spotted problem. On ANY spotted issue, bug, test failure, gate failure, regression, misalignment, inconsistency, unexpected behaviour, or contradiction-between-evidence-and-expectation, the agent MUST activate the structured root-cause-investigation skill

— `superpowers:systematic-debugging` on Claude Code, or the platform-equivalent structured-debugging discipline on any other runtime — **BEFORE proposing, writing, or applying any fix**. This is the **Iron Law: NO FIXES WITHOUT ROOT CAUSE INVESTIGATION FIRST**. The investigation **MUST** run its full four-phase arc: (1) **root-cause phase** — read the actual error / failing assertion / divergent evidence, reproduce deterministically, gather the facts the failure exposes; (2) **pattern phase** — identify what class of defect this is, search for the same pattern elsewhere in the codebase; (3) **hypothesis phase** — form a falsifiable hypothesis about the cause, then prove or disprove it with captured evidence (never `LIKELY` / `probably` / `seems` per §11.4.6); (4) **implementation phase** — only after the cause is a **FACT**, design the fix against the proven root cause. Guess-and-retry, symptom-patching, "let me just try changing this and see", and re-running a failed test hoping it passes ("it's probably transient / flaky") **WITHOUT** a completed root-cause investigation are themselves §11.4.102 violations. Real-world signal this anchor closes: declaring a failure `transient` / `flaky` / `intermittent` / `probably a timing issue` **without captured forensic evidence proving that classification** is simultaneously a §11.4.6 (no-guessing) violation AND a §11.4.7 (demotion-evidence) violation — §11.4.102 makes the corrective response mechanical: the classification is forbidden until the systematic-debugging arc has captured the evidence that proves it.

(B) Mandatory always-loaded skill-discovery discipline (`using-superpowers`). The skill-discovery meta-skill — `superpowers:using-superpowers` on Claude Code, or the platform-equivalent skill-index / capability-discovery discipline on any other runtime — **MUST** be loaded and applied at session start and consulted before any task is begun. Its operative rule: **before acting on ANY request, survey the available skills/capabilities; if ANY skill could apply to the task at hand — even at a 1% chance of relevance — that skill MUST be invoked rather than the agent improvising from memory**. This guarantees the agent reaches for the purpose-built discipline (systematic-debugging for a bug, the brainstorming skill before creative work, the test-driven-development skill before writing a fix, etc.) instead of ad-hoc behaviour. Skipping skill-discovery at session start, or improvising a task that a loaded skill was built to handle, is a §11.4.102 violation.

(C) Mandatory plugin / dependency availability. Every skill plugin, marketplace package, or capability dependency that the project relies on — offered by Claude Code, its skill marketplaces, or the equivalent plugin/extension ecosystem of any other runtime the project uses — **MUST** be installed and loadable on the host

BEFORE the dependent work proceeds. A missing plugin/dependency that blocks a mandated skill (e.g. the `superpowers` plugin absent so `systematic-debugging` cannot launch) is a **release-blocker** until it is installed and confirmed loadable. Install mechanism (generic; runtime supplies the concrete command): the runtime's own plugin/marketplace install path — for Claude Code, the in-session `/plugin` marketplace flow (add the marketplace, install the plugin, confirm the skill appears in the available-skills list); for other runtimes, the documented package/extension installer (npm/pip/cargo/the editor's extension manager). The availability check is itself anti-bluff: a plugin claimed "installed" without the dependent skill actually appearing in the loaded-skills/capabilities list is a §11.4.6 (no-guessing) violation — confirm by observing the skill in the live capability list, not by assuming the install succeeded (install-command exit 0 ≠ skill loadable, mirroring the §11.4.80 `npm exit 0 ≠ working binary` lesson).

(D) AUTOMATIC activation — no operator prompt required (User mandate, 2026-07-13). Verbatim operator mandate (2026-07-13): "Make sure we automatically turn on the systematic-debugging plugin and skills every time we work on any issue! Add this as MANDATORY RULE to our root constitution Submodule and all of its relevant files!" The clause-(A) systematic-debugging activation is AUTOMATIC and the STANDING DEFAULT — the agent MUST auto-activate `superpowers:systematic-debugging` (or the platform-equivalent structured-debugging discipline) on ANY spotted issue / bug / test failure / gate failure / regression / misalignment / inconsistency / unexpected behaviour WITHOUT waiting for an operator prompt, request, or per-issue confirmation. It is engaged the MOMENT an issue is detected — the operator MUST NOT have to ask for it, request it, or re-enable it per issue or per session (composes §11.4.126 default-autonomous-loop, which promotes the standing-default working mode from the first prompt of a session; §11.4.102(D) applies that same always-on-by-default discipline to systematic-debugging activation). Before any such investigation proceeds, the `superpowers` plugin + the `superpowers:systematic-debugging` AND `superpowers:using-superpowers` skills MUST be confirmed installed + loadable per clause (C) — observed in the live capability list, never assumed (§11.4.6 anti-bluff). "I did not activate systematic-debugging because the operator did not ask for it" is itself a §11.4.102(D) violation: clause (A) is NOT a prompt-gated option, it is a standing mechanical default engaged automatically the instant a problem surfaces. §11.4.102(D) STRENGTHENS clause (A) — it makes the

activation trigger prompt-independent and standing-default; it does not weaken it: every clause-(A) four-phase-arc and Iron-Law ("NO FIXES WITHOUT ROOT CAUSE INVESTIGATION FIRST") requirement continues to apply in full.

Composes with §11.4.4 (test-interrupt-on-discovery — §11.4.4 says STOP the cycle the moment a defect is discovered; §11.4.102 says the FIRST action after stopping is to launch systematic-debugging), §11.4.6 (no-guessing — §11.4.102(A) is the procedural enforcement that produces the facts §11.4.6 demands; the forbidden `LIKELY` / `transient` / `flaky` vocabulary is exactly what guess-and-retry emits), §11.4.7 (demotion-evidence — a FAIL → lower-severity demotion requires same-conditions evidence, which is precisely what the systematic-debugging arc captures), §11.4.8 (deep-web-research — runs inside the hypothesis phase), §11.4.43 (TDD-fix — the RED test is written against the root cause the investigation proved), §11.4.70 (subagent-driven — a systematic-debugging investigation is a natural subagent dispatch), §11.4.82(A) (iteration-speedup — already mandates Phase-1 forensic before any speculative patch; §11.4.102 generalises that from "source patch" to ANY fix and adds the always-loaded skill-discovery + plugin-availability invariants), §11.4.92 (multi-pass change-evaluation — the investigation feeds Pass 1 + Pass 4), §11.4.126 (default-autonomous-loop — §11.4.102(D)'s always-on-by-default systematic-debugging activation mirrors §11.4.126's from-first-prompt standing-default working mode).

Classification: universal (§11.4.17) — root-cause-before-fix, skill-discovery-before-task, and dependency-availability are reusable disciplines for ANY project on ANY runtime that exposes a skill/plugin ecosystem; the concrete skill names (`superpowers:systematic-debugging` , `superpowers:using-superpowers`) are the Claude Code instantiation, with the generic "or platform-equivalent" wording carrying the rule to other runtimes per §11.4.81 cross-platform-parity reasoning.

4-layer coverage per §11.4.4(b). Propagation gate `CM-COVENANT-114-102-PROPAGATION` enforces the literal anchor `11.4.102` across the canonical consumer fleet (parent + owned-submodule `CLAUDE.md` / `AGENTS.md` / `QWEN.md`). Paired §1.1 meta-test mutation strips the `11.4.102` literal from a consumer file → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md §11.4.102 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--skip-systematic-debugging` , `--guess-and-retry-OK` , `--symptom-patch-permitted` , `--skip-skill-discovery` , `--plugin-optional` , `--missing-plugin-is-warning` flag exists. The 2026-05-29 user mandate is unambiguous: the moment a problem is spotted the structured root-cause investigation skill activates, `using-superpowers` is always loaded, and every mandated plugin/dependency is installed and loadable — there is no path by which a fix is proposed before the cause is proven.

2026 05 29

§ . . — Continuous parallel-stream working routine (User mandate, - -)

Short tag: `continuous-parallel-stream-routine` .

Forensic anchor — verbatim user mandate (2026-05-29):

"Do this working approach continuously and make it part of regular working routine and add it to the Constitution Submodule documented fully."

This anchor promotes the proven multi-stream operating pattern — repeatedly demonstrated effective across prior autonomous rounds — from an ad-hoc practice into the project's **standing default working routine**. §11.4.87 told the agent to keep working in an endless loop; §11.4.94 told it never to sit idle when something is progressable; §11.4.88 / §11.4.89 detached push and test execution; §11.4.96 catalogued what is safe to do during a long build. §11.4.103 binds those into a single always-on routine and adds the load-bearing invariant the prior anchors only implied: **the main work stream MUST always stay FREE**. It is the conductor's standing operating mode, not a per-request opt-in.

Strengthened (User mandate, 2026-05-29): the parallel-stream floor is raised from ≥ 2 to ≥ 3 and an **auto-backfill rule** is added — the moment any one stream is FULLY done, a new stream MUST immediately start and take its place, so the count NEVER drops below 3 while actionable workable items remain. Verbatim mandate: *"make sure there are always at all times at least 3 parallel working*

streams parallel with main work stream as long there are workable items we can do! as soon as one working stream is fully done, another new MUST start work and take its place."

(A) Main stream stays FREE — ALL commit AND push run detached, never blocking. Every commit operation AND every push operation MUST run in the background (detached via `nohup ... > <log> 2>&1 & + disown`, per §11.4.88's commit-lock-release-immediately + detached-push seam). The conductor's main work stream MUST NOT block on a commit landing or a push completing — the moment the local commit is durable, the main stream returns to the priority queue while the push fans out to every upstream in the background. Blocking the main stream on a push (or a slow mirror) is a §11.4.103 violation (the exact 2026-05-26 §11.4.88 forensic harm — a 5-hour flock hold — generalised to the routine layer).

(B) ≥3 parallel background streams at all times whenever actionable work exists + auto-backfill. Whenever the queue holds three or more non-contending actionable items, the conductor MUST run **at least three** parallel background streams alongside the main stream (subagent-driven per §11.4.70 / §11.4.20, isolated per the §11.4.58 PWU worktree model + §11.4.84 working-tree quiescence). **Auto-backfill rule (User mandate, 2026-05-29):** the moment any one stream is FULLY done (its workable item closed with anti-bluff evidence per sub-rule (E), or merged), a new stream MUST IMMEDIATELY start and take its place by claiming the next-highest-priority non-contending item from the queue — the active-stream count NEVER drops below 3 while actionable workable items remain. Dropping below 3 active streams while three-plus disjoint actionable items wait — or serialising work that could run in parallel — is a §11.4.103 violation. The count may fall below 3 ONLY when there genuinely are no remaining non-contending actionable workable items OR every remaining item is externally blocked (per §11.4.94 / §11.4.97 / §11.4.101 idle-only-when-blocked). The §12.6 60% memory ceiling and the §11.4.58 6-concurrent-agent cap bound the stream count from above (so the standing band is 3–6 background streams).

(C) Most-critical + most-visible first; audio always top. Streams are filled in strict priority order — most-critical and most-user-visible items first — and **audio fixes are always the apex priority** per §11.4.72. When the conductor faces a choice of which item to dispatch onto a free stream, the highest-severity / most-

user-visible / audio item wins; research, refactors, and low-stakes tasks fill remaining capacity only after the critical set is dispatched (per §11.4.42 priority-ordered iteration).

(D) Safe-during-build scope only while a heavy build runs. While the 42 GB containerised AOSP rebuild (§12.9) — or any heavy `gradle` / `m -j` build — is in flight, the parallel background streams MUST restrict themselves to the §11.4.96 SAFE catalogue: investigation / forensic analysis, documentation, test-authoring, gate + paired-mutation authoring, read-only on-device probes, submodule edits, web research, workable-items DB ops, and backgrounded pre-build / meta-test runs. The streams MUST NOT launch a second `gradle` / `m -j` / containerised build concurrent with the running one (§12.8 concurrency-hardening — a second heavy JVM/Soong build storms the host and risks the §12.7 OOM cascade). UNSAFE-during-build operations (§11.4.96 (α)-(θ)) are deferred to between-build windows.

(E) Heavy anti-bluff on every closure. Every item closed by a parallel stream MUST carry real captured evidence and survive the anti-bluff covenant family — root cause proven or explicitly marked `UNCONFIRMED:` / `UNKNOWN:` / `PENDING_FORENSICS:` per §11.4.6, captured-evidence per §11.4.5 + §11.4.69, deterministic consistency per §11.4.50, paired §1.1 meta-test mutations, and the structured root-cause investigation of §11.4.102 the moment any problem surfaces. A stream that closes an item on metadata-only / configuration-only / absence-of-error / grep-without-runtime evidence is a §11.4 PASS-bluff regardless of how parallel the routine looked.

(F) Idle ONLY when genuinely externally blocked. The conductor idles ONLY when every queued item is genuinely blocked on an external result it cannot accelerate (hardware access, network upstream, an in-flight build / test / push completion) OR the operator issues STOP OR §12 host-session-safety demands it (per §11.4.94(A) + §11.4.97). "Nothing visible to do" with progressable non-contending items in the queue is NEVER a valid idle reason. Posing a question and going idle while non-blocked items remain is a §11.4.94 + §11.4.97 + §11.4.103 violation; an unavoidable block parks one work unit, not the whole loop (§11.4.101 maximize-progress-while-blocked).

Composes with §11.4.58 (parallel-development PWU pipeline — §11.4.103 is its standing-routine framing), §11.4.70 / §11.4.20 (subagent-driven default — the parallel streams ARE subagents), §11.4.72 (audio top-priority — sub-rule (C)),

§11.4.87 (endless-loop autonomous work — §11.4.103 is the always-on operating mode that anchor governs), §11.4.88 (background-push — sub-rule (A) is its routine-layer generalisation), §11.4.89 (background test execution — backgrounded tests are one stream class), §11.4.94 (zero-idle priority-first parallel-by-default — sub-rules (B)+(C)+(F) bind it), §11.4.96 (safe-parallel-work-with-long-build catalogue — sub-rule (D) cites it), §11.4.97 (maximum-use-of-idle-time — sub-rule (F)), §11.4.101 (autonomous-decision-over-blocking — maximize-progress-while-blocked), §11.4.102 (systematic-debugging on any spotted problem inside any stream), §11.4.42 (priority-ordered iteration), §11.4.84 (working-tree quiescence for concurrent subagent commits), §12.6 / §12.7 / §12.8 / §12.9 (host-session safety bounds stream count + forbids a second heavy build), §9.2 (data-safety for any destructive op a stream would attempt).

Classification: universal (§11.4.17) — continuous parallel-stream working with a free main stream, priority-first dispatch, and anti-bluff closures is a reusable operating discipline for ANY project on ANY runtime that supports background execution + subagent delegation; nothing in the routine is hardware-, vendor-, or region-specific.

4-layer coverage per §11.4.4(b). Propagation gate `CM-COVENANT-114-103-PROPAGATION` enforces the literal anchor `11.4.103` across the canonical consumer fleet (parent + owned-submodule `CLAUDE.md` / `AGENTS.md` / `QWEN.md`). Paired §1.1 meta-test mutation strips the `11.4.103` literal from a consumer file → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md §11.4.103 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--block-main-stream` , `--synchronous-commit` , `--synchronous-push` , `--single-stream-only` , `--skip-parallel-streams` , `--serialise-actionable-work` , `--idle-without-queue-survey` flag exists. The 2026-05-29 user mandate is unambiguous: continuous parallel-stream working with a free main stream is the regular working routine, not a per-request opt-in.

§ . . — Participant identity, attribution & notification-tagging (User mandate, — —)

Short tag: `participant-attribution-tagging` .

Forensic anchor — verbatim user mandate (2026-05-31):

"Every supported messenger must relate messages to participants (Subscribers/Users); the same logical person may have a different username on every messenger. Workable items must carry who created them and who they are assigned to. Notifications must @-tag the right participant — but never the operator (who drives the system) and never the system agent."

This anchor binds participant identity, workable-item attribution, and notification @-tagging as MANDATORY constraints on every consumer that ships a messenger/ notification surface (Herald and its flavor binaries are the reference implementation; a consuming project and any future messenger-bearing project inherit per §11.4.35). The detailed normative spec is the Herald design contract `docs/design/PARTICIPANT_ATTRIBUTION.md` (model, columns, attribution rules, tagging matrix) — this section restates its load-bearing constraints; implementations code against the contract, not against a paraphrase.

(A) Participant identity — logical person + per-channel handle. Every supported messenger MUST relate inbound + outbound messages to a **Participant** (a logical Subscriber/User). A single logical person/agent MAY carry a DIFFERENT username on every messenger; identity is therefore modelled as a logical subscriber (canonical, messenger-neutral `handle` , `kind` ∈ {human, agent, service}) PLUS per-channel aliases (`channel` , `channel_user_id` , the per-channel `@username` used for tagging). The **canonical handle** is the string stored in attribution columns — the closed set is `claude` (the reserved system-agent sentinel; `kind=agent`) OR a human's canonical handle (a subscriber `@username` , messenger-neutral, resolved per-channel via the alias table).

(B) Operator designation — env var, not a DB flag. The **Operator** is the one human who drives the system via the agent CLI. The Operator is designated by the environment variable `HERALD_<CHANNEL>_OPERATOR_USERNAME` (e.g.

HERALD_TGRAM_OPERATOR_USERNAME , HERALD_SLACK_OPERATOR_USERNAME) — per messenger, NOT a database flag. The Operator is a normal Participant whose canonical handle equals that env value.

(C) Workable-item attribution — created_by + assigned_to . Every workable item MUST carry two canonical-handle columns: created_by and assigned_to . Attribution rules: opened via the agent CLI prompt (operator-driven) → created_by = Operator ; opened by the system/agent detecting an issue/task/improvement/missing-feature → created_by = "Claude" ; received THROUGH the system as a subscriber message → created_by = the sender's resolved canonical @username . assigned_to **defaults to the Operator** and MAY be overridden explicitly (a prompt/message assigning to @someoneelse). Both columns store the canonical handle string and are self-contained in the SSoT + its Markdown export. Legacy items predating this anchor carry empty ("") attribution and MUST still parse + validate.

(D) Notification @-tagging matrix. On any workable-item event, the outbound notification dispatched to each messenger group @-tags the participant(s) who must be aware, resolved to that channel's @username :

```
mentions = {}
if assigned_to is a human handle AND assigned_to !=
Operator:                mentions += assigned_to
if created_by is a human handle AND created_by != Operator AND
created_by != "Claude": mentions += created_by
# "Claude" is NEVER tagged (it is the system agent).
# the Operator is NEVER tagged (drives the system; no self-ping).
# de-dup; for each mention resolve the @username on the target channel —
skip if the participant has no alias there.
```

This satisfies the operator's stated cases exactly: assigned-to-Operator → no tag; opened-by-Operator-assigned-to-another → tag the assignee; opened-by-a-non-Operator-non-Claude-subscriber → tag the creator.

(E) Anti-bluff (MANDATORY, composes with §11.4 / the end-user quality covenant). Every layer ships unit + integration + E2E + full-automation tests producing REAL captured evidence (no metadata-only / absence-of-error PASS, no false positive/negative): real SQLite round-trip with the new columns (byte-identical, including legacy fixtures that DON'T carry the fields); the tagging matrix proven by a truth-table test with a mutation that flips one cell forcing a FAIL; an E2E real-item-

event → real-dispatched-message asserting the exact @username s AND a NEGATIVE case proving the Operator is NOT tagged. Evidence committed under docs/qa/<run-id>/ .

Composes with §11.4 + §11.4.1..§11.4.16 (end-user quality / anti-bluff covenant — sub-rule (E) is bound by it), §11.4.5 / §11.4.69 (captured evidence), §11.4.50 (deterministic consistency), §11.4.91 (workable-item field clarity), §11.4.93 / §11.4.95 (the SQLite workable-items SSoT the created_by / assigned_to columns extend), §1.1 (paired-mutation proof of the tagging matrix).

Classification: universal (§11.4.17) — participant identity, item attribution, and notification @-tagging is a reusable discipline for ANY project that ships a messenger/notification surface with a workable-items registry; the env-var operator designation + canonical-handle model + tagging matrix are vendor-neutral. Projects with NO messenger surface inherit the anchor latently (it binds the moment they ship one) — the §11.4.96 "principle binds even absent the surface" restatement pattern.

4-layer coverage per §11.4.4(b). Propagation gate CM-COVENANT-114-104-PROPAGATION enforces the literal anchor 11.4.104 across the canonical consumer fleet (parent + owned-submodule CLAUDE.md / AGENTS.md / QWEN.md). Paired §1.1 meta-test mutation strips the 11.4.104 literal from a consumer file → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md §11.4.104 in the HelixConstitution submodule; detailed spec Herald docs/design/PARTICIPANT_ATTRIBUTION.md . All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker. No escape hatch — no --skip-attribution , --no-participant-tagging , --tag-operator-anyway , --attribution-later flag exists.

2026 05 31

§ . . — Natural-language intent recognition & clarification (User mandate, - -)

Short tag: intent-recognition-clarification .

Forensic anchor — verbatim user mandate (2026-05-31):

"Users must NOT need to know command syntax (no `COMMAND: ...` prefix). They send a clear natural-language message; the System determines the intent. The System recognizes the commands it has; if none matches it infers the exact intent; if it is totally unable it replies, tags the user (`@user ...`), and asks to clarify precisely. We MUST always do our best to determine exact intent so we never annoy end users. This is a CORE part of the System."

This anchor binds natural-language intent recognition as a MANDATORY constraint on every consumer that ships a messenger/notification-driven command surface (Herald and its flavor binaries are the reference implementation; a consuming project and any future messenger-bearing project inherit per §11.4.35). The detailed normative spec is the Herald design contract `docs/design/INTENT_RECOGNITION.md` (three-tier resolution, command set, the `clarify` action, envelope wiring) — this section restates its load-bearing constraints; implementations code against the contract, not against a paraphrase.

(A) No required command syntax. Users MUST NOT be required to know any command syntax — no `COMMAND: ...` prefix, no rigid grammar. They send a clear natural-language message; the System is responsible for determining the intent. Any design that forces an end user to memorize syntax to be understood is a violation of this anchor.

(B) Three-tier intent resolution (the mandatory discipline). Every inbound subscriber message MUST be resolved to exactly one action via three tiers, in order — the first that succeeds wins:

TIER 1 command recognition – recognize the System's existing command set from natural language and map it to a structured action (a confident, deterministic match → that action).

TIER 2 intent inference – when NO command matches, infer the exact intent from the message (the LLM dispatch maps natural language to the right action) – NEVER guessing.

TIER 3 clarify (fallback) – when neither a command nor a confident intent can be determined, REPLY to the original message, TAG the sender (`@username`, resolved via the §11.4.104 IdentityResolver) and ask a PRECISE clarifying question that NAMES the candidate intents. No guessing, no silent drop.

Tier 3 is the anti-annoyance guarantee: the user is never ignored and never has to learn syntax — at worst they receive a friendly, specific "@user, did you mean X or Y?".

(C) Never guess, never drop. The System MUST NEVER guess an action — a wrong action is worse than a clarifying question (composes with §11.4.6 no-guessing). The System MUST NEVER silently drop or ignore an inbound message. We always do our best to determine the exact intent so end users are never annoyed; only genuine ambiguity reaches Tier 3, and Tier 3 always replies-tags-and-asks rather than failing silently.

(D) Anti-bluff (MANDATORY, composes with §11.4 / the end-user quality covenant). Every tier ships unit + integration + E2E + full-automation tests producing REAL captured evidence (no metadata-only / absence-of-error PASS, no false positive/negative): Tier 1 proven by a truth-table of natural-language messages → expected action+fields PLUS the conservative negatives that MUST fall through to "no match"; Tier 3 proven by an E2E ambiguous-message dispatch whose recorded reply body is EXACTLY @<sender> <specific question> (the

user is tagged + asked, not ignored) AND a NEGATIVE proving a clear command does NOT trigger clarify; a paired §1.1 mutation that breaks the recognizer's confidence guard (forcing a false-match) OR drops the clarify tag MUST FAIL a test. Evidence committed under `docs/qa/<run-id>/` .

Composes with §11.4 + §11.4.1..§11.4.16 (end-user quality / anti-bluff covenant — sub-rule (D) is bound by it), §11.4.6 (no-guessing — a wrong action is worse than a clarifying question), §11.4.104 (participant identity — the clarify reply tags the sender via the IdentityResolver), §11.4.5 / §11.4.69 (captured evidence), §11.4.98 (full-automation — every tier's test is self-driving), §1.1 (paired-mutation proof of the confidence guard + clarify tag).

Classification: universal (§11.4.17) — natural-language intent recognition with a three-tier recognize → infer → clarify discipline is a reusable interaction contract for ANY project that ships a messenger/command surface; the no-syntax mandate, never-guess rule, and reply-tag-and-ask fallback are vendor-neutral. Projects with NO messenger surface inherit the anchor latently (it binds the moment they ship one) — the §11.4.96 "principle binds even absent the surface" restatement pattern.

4-layer coverage per §11.4.4(b). Propagation gate `CM-COVENANT-114-105-PROPAGATION` enforces the literal anchor `11.4.105` across the canonical consumer fleet (parent + owned-submodule `CLAUDE.md` / `AGENTS.md` / `QWEN.md`). Paired §1.1 meta-test mutation strips the `11.4.105` literal from a consumer file → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md §11.4.105 in the HelixConstitution submodule; detailed spec Herald `docs/design/INTENT_RECOGNITION.md` . All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker. No escape hatch — no `--require-command-syntax` , `--guess-intent-ok` , `--skip-clarify` , `--drop-on-ambiguous` flag exists.

§ . . — Docs Chain — mechanical documentation/DB sync engine (Operator mandate, — —)

Short tag: `docs-chain-mechanical-sync` .

Forensic anchor — operator mandate (2026-05-31):

Docs Chain is the canonical mechanical enforcer of the documentation-sync mandates. Consuming projects **MUST** use the engine instead of ad-hoc per-project doc-sync scripts; they register their own chains via per-context YAML; the engine never fakes a transform and every `sync / verify` carries real captured evidence.

Docs Chain — the `vasic-digital/docs_chain` engine (a universal, Go-implemented, **bidirectional document-and-database dependency-propagation engine**) — is the canonical mechanical enforcer of the family of documentation-sync mandates this Constitution already declares. This is a **MANDATORY** constraint binding every consuming project. The detailed normative spec is the engine's own documentation set — `docs/CONSTITUTION_INTEGRATION.md` (distribution + inheritance + the anchor mapping table) and `docs/USE_CASE_CATALOGUE.md` (the ready-to-use chain recipes); this section restates its load-bearing constraints, it does not redefine them. Implementations code against the engine's contract, not against a paraphrase.

(A) Use the engine, never ad-hoc scripts. Consuming projects **MUST** drive their documentation/DB sync through Docs Chain — consumed **by reference** via the flat-layout sibling (the `~/Projects/docs_chain` working tree during local development / the path the constitution exposes once the Phase-6 submodule pointer lands), inherited the same way §11.4.80 inherits the `codegraph_*` scripts: referenced at the constitution-exposed path, **NEVER copied** into the consumer. Ad-hoc per-project doc-sync scripts (`sync_issues_docs.sh` , `generate_*_summary.sh` , `sync_*_export.sh` , `update_readme_doc_links.sh` , et al.) are superseded by the engine and **MUST** be retired as each matching context is registered.

(B) Consumer-owned contexts via config. The engine is project-agnostic; it carries NO project-specific path or chain. The consumer registers its chains as data via `.docs_chain/contexts/*.yaml` at its own project root — one independent context per file — per the §11.4.28 decoupling rule (the submodule never carries project-specific context; the consumer injects its chains via config). `state.json` (at `<project-root>/docs_chain/state.json`, regenerated on demand by `docs_chain sync`) and the `*.docs_chain.tmp` staging sidecars are **gitignored**; the consumer owns those `.gitignore` entries.

(C) The sync anchors it mechanizes. Docs Chain is the mechanical engine behind a family of existing documentation-sync anchors; once a consumer registers the matching context it enforces the anchor **in place of the ad-hoc script** — see the `docs/CONSTITUTION_INTEGRATION.md` mapping table: §11.4.12 (Issues_Summary always-sync), §11.4.53 (Fixed_Summary parity), §11.4.45 (Status.md maintenance), §11.4.56 (Status_Summary two-audience parity), §11.4.57 (README doc-link section), §11.4.59 (README always-sync export), §11.4.60 (documentation composite-sync), §11.4.65 (universal markdown export), §11.4.86 (roster/corpus auto-sync), §11.4.93 (workable-items DB single source of truth), §11.4.95 (DB tracked + WAL-checkpoint commit), §12.10 (CONTINUATION maintenance), and §11.4.44 (document revision header — exports kept in sync). Engine-wide guarantees: change detection by **content hash, NOT mtime** (§11.4.86); **atomic-rename + SQLite-transaction commit** with rollback on any transform error and optional hardlink backup (§9.2 zero-risk data safety); a both-dirty `sync` pair surfaces a **conflict — never a silent merge** (exit 2, §11.4.6 no-guessing); `verify` is the deterministic CI / pre-build sink-side gate over byte-stable transforms (§11.4.50); per-run captured evidence lands at `qa-results/docs_chain/<run-id>/` (§11.4.69).

(D) It does NOT replace authoring discipline. Docs Chain replaces only the **mechanical sync** an anchor requires — it does NOT replace the authoring discipline of any anchor. The source author still writes the §11.4.44 revision header into the source; the engine then keeps every derived export in sync. A consumer must not treat the engine as licence to skip authoring.

(E) Anti-bluff (MANDATORY, composes with §11.4 / the end-user quality covenant). Docs Chain NEVER fakes a transform: a missing `pandoc` / `weasyprint` (or any absent tool) surfaces a typed `ToolAbsentError` (`IsToolAbsent`) and an honest §11.4.3 SKIP-with-reason — **never a fake PASS**,

never a partial write. Every `sync / verify` run carries real captured evidence (the `qa-results/docs_chain/<run-id>/` artefact). A metadata-only / "absence-of-error" / config-only PASS at the sync layer is a §11.4 PASS-bluff.

Status note (§11.4.6 no-guessing). The engine's Phases 1–3 (core DAG + content-hash + adapters + atomicity/orchestrator) AND the `cmd/ CLI / per-context` YAML loader (Phase 4) are IMPLEMENTED + tested —

`cmd/docs_chain/main.go` registers the `sync / verify / diff` (alias of `verify`) / `doctor` subcommands wiring the config loader → state → orchestrator, and `go test -count=1 ./...` is GREEN at 7/7 packages (incl. `internal/orchestrator spec_wiring_test`, both PASS). FACT (§11.4.6): this CLI/loader lives in the working tree but is UNTRACKED (`git ls-files docs_chain/cmd/docs_chain/main.go` is empty) and the engine is NOT yet a registered git submodule (`git submodule status` shows no `docs_chain` entry; in-tree at `./docs_chain` , HEAD = parent HEAD). The constitution-submodule distribution (Phase 6) remains PLANNED + OPERATOR-GATED. This anchor states the binding contract; consumers wire the engine as each phase lands, and MUST NOT claim working behaviour a phase has not yet shipped.

Classification: universal (§11.4.17) — a content-hash bidirectional doc/DB propagation engine is a vendor-neutral mechanism reusable by ANY governed project. Projects with no derived-export or DB-sync surface inherit the anchor latently (it binds the moment they ship one) — the §11.4.96 "the principle binds even absent the surface" restatement pattern.

4-layer coverage per §11.4.4(b). Propagation gate `CM-COVENANT-114-106-PROPAGATION` enforces the literal anchor `11.4.106` across the canonical consumer fleet (parent + owned-submodule `CLAUDE.md / AGENTS.md / QWEN.md`). Paired §1.1 meta-test mutation strips the `11.4.106` literal from a consumer file → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Composes with §11.4 + §11.4.1..§11.4.16 (end-user quality / anti-bluff — sub-rule (E) is bound by it), §11.4.6 (no-guessing — both-dirty `sync` → conflict, never silent merge), §11.4.12 / §11.4.45 / §11.4.53 / §11.4.56 / §11.4.57 / §11.4.59 / §11.4.60 / §11.4.65 / §11.4.86 / §11.4.93 / §11.4.95 / §12.10 / §11.4.44 (the documentation-sync anchors it mechanizes), §11.4.28 (engine/context decoupling), §11.4.80

(inherited-by-reference, never-copied — the codegraph pattern), §9.2 (atomic commit + rollback), §11.4.50 (deterministic `verify` gate), §11.4.69 / §11.4.5 (captured evidence), §1.1 (paired-mutation propagation proof).

(F) Write-seam hook enforcement + real-time-sync honest boundary

(extension, 2026-07-15; operator mandate — the Continuation / memory /

knowledge artefacts MUST ALWAYS be in sync). The documentation/DB sync this engine mechanizes MUST be ENFORCED AT THE WRITE SEAMS, never left to agent vigilance: the consuming project MUST wire the engine's `verify` (and, where safe, `sync`) into (1) the COMMIT seam — a pre-commit / commit-wrapper hook that REFUSES a commit whose staged set touches a chain SOURCE while its derived exports / DB rows are stale; (2) the BUILD seam — the pre-build gate; and (3) the CONSTITUTION-PULL seam — the §11.4.164 `post_update_hook.sh` invocation. A stale `CONTINUATION.md` (§12.10), a stale standing session-resumption file (§11.4.131), or a stale memory / knowledge artefact while the live state has ALREADY moved is a violation of the SAME severity class as a stale export — the resumption guarantee silently breaks and a fresh session resumes on fiction. Honest boundary (§11.4.6): hook-wired seam enforcement makes the sync EVENTUALLY-CONSISTENT-AT-EVERY-WRITE — it is NOT literal continuous real-time; genuine continuous real-time sync requires a filesystem-watch daemon (inotify / FSEvents watcher running the chain on change). A project that CLAIMS "real-time sync" while shipping only seam hooks is bluffing the mechanism: state the seam-based guarantee honestly, and carry the watch daemon as a TRACKED upgrade path (§11.4.197 — driven to completion or explicitly closed, never left unwired).

Canonical authority: this Constitution.md §11.4.106 in the HelixConstitution submodule; detailed spec the Docs Chain engine docs `docs/`

`CONSTITUTION_INTEGRATION.md` + `docs/USE_CASE_CATALOGUE.md` (`vasic-digital/docs_chain`). All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker. No escape hatch — no `--ad-hoc-sync-ok`, `--skip-docs-chain`, `--fake-transform`, `--sync-evidence-optional` flag exists.

§ . . — Anti-bluff AV/test-validation techniques mandate (User-driven research, — —)

Short tag: `av-liveness-oracles` .

Forensic anchor — captured incidents (2026-06-02, genericised): a test asserting "media is playing on the target output" PASSEd on the strength of a SINGLE captured frame that showed "a picture" — but the picture was a FROZEN / STALE frame left over from the previously-played content (a stale-producer / stuck-decoder failure mode), so the feature was broken for the end user while the test was green. A sibling incident: media briefly FLASHED on the WRONG output (the source/primary surface) for ~1 s before routing to the intended output, and no test sampled the wrong output during the transition window, so the glitch shipped undetected. A third class: a comparator (color / freeze / OCR oracle) that was itself broken — it PASSEd on a deliberately-degraded artefact — meaning the analyzer, not just the feature, was the bluff. §11.4.5 mandates *captured* evidence and a baseline presence pass; §11.4.107 raises the bar to **liveness + correct-routing + self-validated-analyzer** so a single still frame, a wrong-output flash, or a broken oracle can no longer pass.

The mandate. Every test asserting that audio or video output is genuinely playing / advancing for the end user MUST satisfy ALL of the following — each is a §11.4 PASS-bluff to omit:

1. **A single captured frame is NOT proof of playback.** Prove LIVE, ADVANCING frames over a steady-state window ($\geq N$ frames across $\geq$ a few seconds) via a **freeze-detection oracle** — a near-duplicate detector (the platform's frame-difference / `freezedetect` -class filter OR a perceptual-hash adjacent-frame distance), **NOT a byte-identical compare** (a re-encoded / 1-pixel-noisy stale frame defeats byte-identity; keep byte-identity only as a zero-cost early-out pre-filter). Adjacent frames MUST differ above a calibrated threshold; $\geq$ a sustained interval of near-identical frames $\Rightarrow$ FROZEN $\Rightarrow$ FAIL.
2. **An independent frame-advance counter from the platform's compositor / decoder telemetry** MUST increase across the window — an oracle in a DIFFERENT domain from the pixel check (compositor frame-timestamp advance, decoder `framesDecoded` / `framesRendered` , or equivalent). A flat counter over the window $\Rightarrow$ stuck decoder $\Rightarrow$ FAIL, even if the pixels appear to

move. Keeping a pixel-domain oracle AND a telemetry-domain oracle is metamorphic redundancy: a coordinated false-PASS now requires two independent failures.

3. **Loading / buffering is a DISTINCT state — wait for genuine playback-start before judging liveness.** Never false-FAIL a still-loading stream, never false-PASS a loading spinner. Poll the decoder frame-advance counter / first-advancing-frame / buffering flag until playback genuinely starts (generous per-class timeout). Timeout + content endpoint unreachable ⇒ SKIP-with-reason per §11.4.3 (`network_unreachable_external` / `geo_restricted`); timeout + reachable ⇒ FAIL with captured evidence. Only AFTER confirmed playback-start run the liveness checks over the steady-state window.
4. **Not-stale-from-previous cross-check.** Before asserting new content N is live, confirm its first frame DIFFERS from the previous content's last captured frame (perceptual-hash mismatch). Identical ⇒ the "new" content is a leftover frozen frame ⇒ FAIL. Critical for switch-between-items sequences.
5. **Measured FPS / no-lost-frames within tolerance** of the source's expected rate (where the source rate is knowable) or a sane floor otherwise; drop ratio above budget ⇒ FAIL (§11.4.5 frame-health).
6. **No-flash-on-the-wrong-output during a routing transition.** When content is meant to appear ONLY on a target output, sample the NON-target output at high frequency (e.g. every ~200 ms) across the transition window; ANY sampled wrong-output frame containing the content ⇒ FAIL (brief-flash glitch). (Where a content-protection regime legitimately forces a particular output, that regime is the expected case, not a fail — classify it explicitly, never by guessing per §11.4.6.)
7. **Drive tests through the realistic user path (feed / UI), not deep-link shortcuts.** Routing / transition / alternation bugs manifest on the path a real user takes (selecting from a feed / list via the UI); deep-link / direct-dispatch shortcuts bypass the very transition paths where these bugs live. Where the UI is not introspectable, fall back AND flag operator-attended per §11.4.52 — never fake-PASS. (This is the §11.4.48 UI-driven discipline applied to the liveness battery.)

8. **The oracle problem — use metamorphic relations when there is NO golden source.** For content the project does not own (streamed media), there is no golden reference, so full-reference comparison is impossible. Encode metamorphic relations instead: the SAME content captured on output-A vs output-B must be structurally equivalent (after resolution normalisation); a paused stream's frame-advance counter must STOP; a 2× speed must roughly double the frame-advance rate; the same clip played twice must produce equivalent liveness metrics. Each is a relation, not a fixed expected value — robust to content variety.
9. **Full-reference quality metrics vs a golden source for content the project DOES own.** Where a source reference exists (local-file playback of owned media, a test corpus), assert perceptual-quality / color fidelity (SSIM / VMAF / ΔE_{2000} -class) against the temporally-aligned source so a pale / desaturated / wrong-range / soft-scaled render is caught across motion, not just at one sampled frame.
10. **Mutation-test every analyzer with a golden-good + golden-bad fixture pair** so the analyzer itself provably cannot bluff. For every comparator (freeze / motion / color / sharpness / aspect-ratio / FPS / OCR / audio-presence) keep a clean control fixture (must PASS) AND a deliberately-degraded fixture (must FAIL — frozen clip, black-gap clip, desaturated clip, downsampled clip, stretched clip, half-rate clip, overlay-baked frame, silent track). An analyzer that PASSES its golden-bad fixture is a bluff gate. This is the §1.1 paired-mutation discipline applied to the AV *analyzers*, not only the grep gates. Document any behaviourally-equivalent mutant expected NOT to flip so a survivor there is not a false alarm.
11. **Per-channel audio integrity, not a single aggregate number.** A single aggregate RMS misses a dead channel in a multichannel stream and a near-silent / DC-offset stream that scrapes past a bare floor. Assert per-channel RMS / loudness (an EBU-R128-class integrated/momentary loudness + per-channel level + silence-region census + DC-offset / NaN-Inf check); EVERY declared channel above floor. On-device, take an underrun / XRUN census from the audio subsystem — a starved pipeline is a defect even if the captured waveform "has sound".

12. **OCR overlay / subtitle detection needs a per-word confidence floor + a region-of-interest** to avoid BOTH a false-positive (noise read as a hostile overlay ⇒ false-FAIL of a working feature) AND a false-negative (a real overlay missed ⇒ false-PASS). Discard sub-confidence words; OCR the ROI where the text lives (bottom-third for subtitles, centre for dialogs); emit per-finding confidence so a reviewer sees WHY a finding fired. Where a subtitle must appear on one output and NOT another, assert both sides (a property relation per point 8).

13. **Thresholds MUST be calibrated on the project's own fixtures, not hardcoded from literature (§11.4.6 no-guessing).** Every freeze / motion / SSIM / VMAF / loudness / OCR-confidence threshold is recorded as captured evidence derived from a known-good and known-bad sample on the project's own capture path — never a literature constant asserted as fact. The capture path's own re-encode loss is part of the baseline.

Honest gaps (§11.4.6). True physically-displayed (photon) frame rate has no clean software oracle — the compositor / decoder counters measure the *presentation pipeline*, not photons at the sink; a high-speed external camera rig is the only ground truth. The counters are sufficient to catch a stuck / half-rate decoder and MUST be used, with the photon-FPS gap flagged honestly rather than claimed. Broadcast-QC vendor tooling validates that freeze / duplicate / lost-frame / black detection is industry-standard practice, but the exact proprietary thresholds are undocumented; the project adopts the published open-method analogues and calibrates per point 13.

4-layer coverage per §11.4.4(b). Per-fix: pre-build gate (e.g. a `CM-AV-LIVENESS-NO-FROZEN-FRAME` -class gate asserting every test that claims output-is-playing references the liveness battery — freeze oracle + frame-advance counter + not-stale cross-check + FPS — not a single-frame check; plus an analyzer-self-validation gate wiring the golden-good/golden-bad fixtures into the project's meta-test) + on-device / runtime test + **paired §1.1 meta-test mutation** (a test asserting output-is-playing with ONLY a single-frame check → gate FAILs; an analyzer that PASSes its golden-bad fixture → self-validation FAILs) + HelixQA Challenge entry for the user-visible feature. Every such PASS is emitted via `ab_pass_with_evidence` citing the motion / not-stale / frame-advance / fps / per-

channel-loudness / metamorphic / self-validation artefacts (feature classes `video_display` , `audio_output` , `subtitle_render` per §11.4.69) — NEVER a single screenshot.

Classification: universal (§11.4.17) — freeze-detection oracles, frame-advance telemetry, loading-aware playback-start waits, not-stale cross-checks, metamorphic relations, full-reference quality metrics, golden-good/golden-bad analyzer self-validation, per-channel loudness, and confidence-floored OCR are platform-neutral AV / test-validation techniques reusable by ANY project that validates media playback or any pixel/audio output; the project supplies the concrete capture mechanism (its compositor / decoder telemetry, its sink-introspection API per §11.4.13, its capture tool) and the calibrated thresholds per §11.4.35.

Composes with §11.4.5 (strict expansion of its video + audio quality bar to liveness + correct-routing), §11.4.6 (no-guessing — thresholds calibrated + captured, loading-vs-frozen disambiguated by evidence, content-protection regime classified not guessed), §11.4.50 (deterministic-consistency — the liveness battery runs N iterations identically), §11.4.68 (positive sink-side / downstream evidence — the frame-advance counter / sink-introspection cross-check is its video analogue), §11.4.69 (universal sink-side positive-evidence taxonomy — `video_display` / `audio_output` / `subtitle_render` classes), §11.4.85 (stress + chaos — liveness under sustained load + transition chaos), plus §11.4.2 (recorded evidence), §11.4.3 (topology SKIP-with-reason for unreachable content), §11.4.13 (sink-side introspection as one oracle), §11.4.48 (UI-driven path), §11.4.52 (autonomous validation; operator-attended fallback only when UI not introspectable), §1.1 (paired-mutation discipline applied to the analyzers).

Propagation. Propagation gate `CM-COVENANT-114-107-PROPAGATION` enforces the literal anchor `11.4.107` across the consumer fleet (parent + owned-submodule CLAUDE.md / AGENTS.md / QWEN.md); paired §1.1 meta-test mutation strips the `11.4.107` literal from a consumer file → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md §11.4.107 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker. No escape hatch — no `--single-frame-proves-playback` , `--skip-liveness` , `--byte-identical-freeze-OK` , `--no-frame-advance-counter` , `--skip-not-stale-check` , `--allow-wrong-output-flash` , `--deep-link-shortcut-OK` , `--unvalidated-analyzer-OK` , `--aggregate-rms-suffices` , `--hardcoded-thresholds-OK` flag exists.

4 5

2026 06 03

§ . . – Four-layer fix-verification
+ runtime-signature-as-definition-of-done
mandate (systematic-debugging Phase . ,
– –)

Short tag: `runtime-signature-definition-of-done` .

Forensic anchor (genericised, 2026-06-03). Across one batch, multiple fixes were "green" at every gate yet NOT working for the end user: a change present in the source tree and passed by both the pre-build (source) gate and the post-build gate never reached the boot-time kernel command-line embedded in the deployed boot artifact, so the corresponding output feature was dead despite all-green; sibling fixes correctly built into the system image were masked by stale per-user overlay copies left from a previous deployment (the running code was the stale overlay, not the freshly-built system code), so they merely *looked* broken; deployment did not produce a clean state (the mutable overlay was not wiped), so the stale shadow survived; and the validation suite ran against whatever was running — which could be the stale shadow — so it reported green on code that was never actually exercised. Per `superpowers:systematic-debugging` Phase 4.5: when each fix reveals a fresh "fixed-but-not-working" in a *different* place, that is NOT a series of independent bugs — it is ONE architectural VERIFICATION flaw, and patching each symptom is thrashing.

The architectural flaw (FACT, traced). A fix must cross FOUR distinct layers, and "fixed" at one layer does NOT imply fixed at the next: (1) **SOURCE** — the change is committed in the source file (what a grep-the-source pre-build gate checks); (2) **ARTIFACT** — the change's BYTES actually landed in the produced build artifact (image / bundle / installer / boot artifact / embedded command-line); (3) **RUNTIME-ON-CLEAN-TARGET** — the change is active on a *clean* / *fresh* deployment, the layer the end user experiences, with no stale overlay shadowing the deployed

code; (4) **USER-VISIBLE** — the feature actually works for the end user (the §11.4.5 / §11.4.69 captured-evidence layer). Green at layer 1 is the cheapest, least conclusive signal; nothing being green at layer 1 (or even layer 2) tells you about layers 3–4. A gate that verifies ONLY the source layer is insufficient and is itself a §11.4 bluff surface.

The mandate (ALL must hold).

1. **Runtime-signature-as-definition-of-done.** A fix is DONE only when its declared **runtime signature** is verified on a CLEAN / fresh deployment (the layer the end user experiences). Source-committed $\neq$ artifact-contains-it $\neq$ active-on-clean-target $\neq$ user-visible-working — these are FOUR distinct claims; proving one does NOT prove the next.
2. **Every fix declares ONE machine-checkable runtime signature** — a single observable on a clean target that proves the fix is BOTH active AND working (a property read from the running system, a downstream/sink-side report per §11.4.13, a captured-evidence assertion per §11.4.5 / §11.4.69, or a counter/state delta — never a re-grep of the source). The **registry of these per-fix runtime signatures is the SINGLE SOURCE OF TRUTH for "fixed"** — it REPLACES "a gate greps the source" as the definition of done.
3. **Verification gates MUST span all four layers**, per fix: **source** (pre-build / pre-merge), **artifact** (post-build — assert the change's BYTES landed in the produced artifact, NOT merely that the source still contains the change; cf. the embedded-command-line / stripped-asset failure mode), **runtime-on-clean-target** (post-deploy — assert the runtime signature on a freshly-deployed clean target), **user-visible** (captured evidence per §11.4.5 / §11.4.69). A gate that asserts only the source layer is insufficient and is itself a §11.4 bluff surface.
4. **Eliminate the stale-deployment / shadow layer — by construction.** Deployment MUST yield a CLEAN state (wipe the mutable overlay / per-user mutable state that can shadow the deployed build) **OR** a pre-validation assertion MUST prove `running-artifact == built-artifact` (no stale override masks the freshly-deployed build) BEFORE any validation runs. Validation that runs against possibly-stale deployed state is INVALID and any PASS it produces is a §11.4 PASS-bluff — the test exercised code that was never actually deployed.

5. **Meta-rule (composes §11.4.102 Phase 4.5).** ≥ 3 "fixed-but-not-working" discoveries within a single cycle are the signal of an architectural VERIFICATION flaw, NOT three independent product bugs. On the 3rd such discovery the agent MUST STOP patching symptoms (per §11.4.4 test-interrupt-on-discovery), fix the VERIFICATION pipeline itself (close whichever of the four layers is unverified), and re-certify EVERY item in the batch through the corrected pipeline on a clean target.
6. **A batch is "validated" only after COMPREHENSIVE per-item runtime-signature verification on a clean baseline** — NOT after the batch-touched items' own tests pass. Spot-validating only the touched items (per §11.4.40's full-suite reasoning) misses the cross-layer gap: a touched item can be source-green while its runtime signature is dead on a clean target.
7. **Machine-created evidence at every layer (extension, 2026-08-15 — Point 8 of the 9-point mandate).** At each of the four verification layers (SOURCE / ARTIFACT / RUNTIME-ON-CLEAN-TARGET / USER-VISIBLE), the layer's PASS MUST cite MACHINE-CREATED, MACHINE-VERIFIABLE evidence PRODUCED AT THAT LAYER by the layer's own gate — a schema-parseable file (static-analysis JSON + review verdict at SOURCE; build log + fingerprint checksum + byte-check at ARTIFACT; on-target verdict + runtime-signature + telemetry at RUNTIME; OCR/vision/audio/sink-probe oracle output at USER-VISIBLE) — cited by path + sha256 + timestamp per §11.4.207. A PASS at a higher layer that cites evidence from a lower layer ALONE is a §11.4.226 wrong-layer violation and does NOT satisfy the higher layer's obligation. Every evidence analyzer is §11.4.107(10) golden-good / golden-bad / negative-control validated; an analyzer that passes its golden-bad fixture voids every verdict it ever emitted (§11.4.201 self-validation). Operator eyeballing, narrative-only PASS, and prediction of success are FORBIDDEN — §11.4.262 binds this at every gate. Propagation: the existing CM-COVENANT-114-108-PROPAGATION stays unchanged (literal 11.4.108 unchanged, stays GREEN); CM-RUNTIME-SIGNATURE-REGISTRY is EXTENDED with the per-layer machine-evidence-file-schema check; a new paired mutation that forges a runtime PASS whose only evidence is a source-grep transcript → the gate FAILS.

Classification: universal (§11.4.17) — the four-layer source → artifact → runtime-on-clean-target → user-visible verification gap, the per-fix runtime-signature registry as the definition of done, the artifact-byte-presence gate, the clean-deployment-or-

equal-artifact pre-validation assertion, and the ≥ 3 -discoveries-means-architectural-flaw meta-rule are platform-neutral verification-pipeline disciplines reusable by ANY project that builds an artifact, deploys it, and validates the deployed result; the consuming project supplies the concrete artifact format, its clean-deployment / equal-artifact mechanism, and its per-fix runtime-signature observables per §11.4.35.

Composes with §11.4.1 (FAIL-bluffs forbidden), §11.4.2 (recorded evidence — the runtime signature IS the recorded evidence), §11.4.4 (test-interrupt-on-discovery + retest-from-clean-baseline — clause 5's STOP-and-fix-the-pipeline is its §11.4.108 specialisation; "clean baseline" is exactly clause 4's clean deployment), §11.4.5 (captured-evidence quality — the user-visible layer), §11.4.6 (no-guessing — every layer asserted by captured evidence, never assumed to have propagated), §11.4.27 (no-fakes-beyond-unit — runtime signatures exercise the real deployed system), §11.4.40 (full-suite retest — §11.4.108 adds the cross-layer per-item runtime-signature dimension clause 6 demands), §11.4.46 (validate-recent-work-before-full-suite — the clean-baseline / equal-artifact pre-flight is its §11.4.108 pre-condition), §11.4.50 (deterministic consistency — the runtime signature verifies identically across N iterations), §11.4.52 (autonomous-validation — runtime signatures are machine-checkable without an operator), §11.4.69 (universal sink-side positive-evidence taxonomy — a runtime signature is a taxonomy-class observable), §11.4.102 (systematic-debugging — Phase 4.5 architectural-flaw recognition IS clause 5's trigger).

Propagation. Propagation gate `CM-COVENANT-114-108-PROPAGATION` enforces the literal anchor `11.4.108` across the consumer fleet (parent + owned-submodule `CLAUDE.md` / `AGENTS.md` / `QWEN.md`); paired §1.1 meta-test mutation strips the `11.4.108` literal from a consumer file → the gate FAILs. A recommended per-fix gate `CM-RUNTIME-SIGNATURE-REGISTRY` asserts every fix declares a runtime signature verified at the artifact + runtime-on-clean-target layers (not only the source layer); paired §1.1 mutation downgrades a fix's verification to source-only → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md §11.4.108 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker. No escape hatch — no `--source-green-is-done` , `--skip-artifact-byte-check` , `--validate-against-running-state` , `--no-clean-deployment` , `--skip-runtime-signature` , `--spot-validate-touched-only` flag exists.

\$. . — Mandatory Anti-Forgetting Enforcement: PreToolUse Guard Hook + Subagent Constitutional Preamble + Orchestrator Pre-Action Checklist (Operator mandate)

Short tag: `anti-forgetting-enforcement` .

Forensic anchor — operator mandate:

UNCONFIRMED: pending operator's verbatim anti-forgetting mandate quote (tracked: PENDING item 3 in the current session's CONTINUATION.md handoff).

Background context (UNCONFIRMED quote omitted; factual description only): the motivating incident is documented in `docs/AGENT_GUARDRAILS.md` and in `scripts/hooks/guard-forbidden-commands.sh` 's in-source header. During an on-device-API build, emulator subagents ran raw host-direct `emulator / adb` commands instead of routing through the Containers submodule as required by the governing clauses (§6.X / §6.V / §6.AG in the consuming-project layer). The root cause was that the orchestrator **forgot to inject the rule into subagent prompts**. A rule the orchestrator forgets to paste is not enforcement — it is a social contract with the agent's memory, and that contract is broken by every cold-session start, context-window exhaustion, and sub-agent dispatch. The mechanical fix is twofold: (1) a **PreToolUse guard hook** that blocks the forbidden command classes at the tool-call boundary unconditionally, and (2) a **canonical subagent preamble document** that the orchestrator pastes verbatim into every dispatch, plus a **pre-action checklist** for the orchestrator's own actions. Together these are the "anti-forgetting" enforcement layer: the hook is the floor (commands cannot be executed even if the agent forgot every rule), the preamble is the ceiling (the full live ruleset, not just what the hook can pattern-match). This two-layer model was introduced in

a consuming project (commit family on `master` , 2026-06-01/02) and is explicitly designed to be universal — `scripts/hooks/guard-forbidden-commands.sh` carries `Classification: universal` in its header and is portable by design (no project-specific paths, no `jq` hard dependency).

Rule. Every HelixConstitution-consuming project MUST deploy the §11.4.109 anti-forgetting enforcement layer consisting of three components, all maintained together:

(A) PreToolUse guard hook — the mechanical floor.

A shell script MUST be wired as a `PreToolUse` hook in the AI agent runtime's project-scoped settings file (e.g. `.claude/settings.json` for Claude Code) to intercept every `Bash` tool call before execution and BLOCK the forbidden command classes at the tool-call boundary. The hook:

1. Reads the tool invocation as JSON from stdin (`.tool_name` + `.tool_input.command`).
2. Exits 0 for non-Bash tools (pass-through) and empty/missing commands.
3. For Bash tool calls, pattern-matches against the following mandatory blocked classes, each mapped to its governing clause:

Blocked class	Governing clause(s)	Examples
Raw host-direct Android emulator launch / APK in- stall / instrument- ation run	§11.4.76 (Containers mandate) + any project- layer emulator clause (e.g. §6.X / §6.V / §6.AG in a consuming project)	<code>emulator -avd ... , \$AN- DROID_HOME/emulator/ emulator ... , adb install ... , adb -s ... install ... , am instrument ...</code>
Force-push / hook bypass / signature bypass	§11.4.41 (pre-force-push merge-first) + §11.4.71 (pre-push fetch) + pro- ject-layer §6.T.3 equival- ent	<code>git push --force , git push -f , git push --force-with- lease , --no-verify , --no- gpg-sign</code>
Privilege escala- tion	§11.4 no-sudo (project- layer §6.U equivalent)	<code>sudo ... , su , su - , su -l ...</code>
Host power man- agement	§12 host-session-safety (all consuming projects)	<code>systemctl suspend/hibern- ate/poweroff/reboot/halt/ kill-user/kill-session , loginctl ... , pm-suspend/ hibernate , shutdown ...</code>

4. Exits 2 (block) when a match is found; the stderr text printed is fed back to the agent as the refusal reason, citing the governing clause.
5. Supports a documented-exception escape hatch: a command containing the literal marker `# guardrails:allow <reason>` is WARNED (stderr) but NOT blocked — EXCEPT for host-power commands, which are categorically non-overridable.

The **canonical implementation** of this hook lives at `constitution/scripts/hooks/guard-forbidden-commands.sh` in this constitution submodule. Consuming projects MUST reference it at that path (inherited by reference per §11.4.80) — NEVER copy it locally (a copy diverges silently). The script is portable: no `jq` hard dependency (built-in `awk` fallback), no consumer-project-specific paths, no hardcoded credentials.

Anti-bluff requirement: a hermetic test suite (at minimum 20 cases) MUST verify every blocked class exits 2, every allowed command exits 0, the escape hatch fires for non-power classes, and the host-power class rejects even with the escape marker. Paired §1.1 mutation: remove the `emulator -avd` pattern from the hook → the emulator-gate case exits 0 → test FAILs → restore → test PASSes.

(B) Subagent Constitutional Preamble — the semantic ceiling.

A canonical document MUST exist at `docs/AGENT_GUARDRAILS.md` (or the project's equivalent governance preamble path, recorded in the project's `CLAUDE.md`) containing:

1. **The SUBAGENT CONSTITUTIONAL PREAMBLE block** — a verbatim, self-contained text block covering the top-10 mandatory constraints that cannot be pattern-matched by the hook alone (anti-bluff intent, resource caps, evidence honesty, continuation maintenance, hardcoding prohibition, distribute gate, remote policy, no-guessing vocabulary, etc.) plus the constraint classes that the hook already enforces (emulators, force-push, sudo, host-power) for completeness.
2. **Clear instruction that the orchestrator MUST paste this block verbatim into every subagent dispatch.**

The preamble is the semantic complement to the hook: the hook catches command-class violations; the preamble pre-informs the agent of all rules it must follow before it even issues a command. Together they eliminate the "agent forgot because the orchestrator didn't paste it" failure class.

Anti-bluff requirement: the preamble MUST be structured so it can be mechanically verified to be present in `docs/AGENT_GUARDRAILS.md` by `check-constitution.sh` (or the project-equivalent sweep); the document MUST contain the anchor literal `11.4.109` once this clause lands. Paired §1.1 mutation: remove the preamble's hook-reference sentence → constitution checker detects absence → FAIL.

(C) Orchestrator Pre-Action Checklist — guard against self-forgetting.

The same `docs/AGENT_GUARDRAILS.md` document MUST contain an **ORCHESTRATOR PRE-ACTION CHECKLIST** covering at minimum:

- Before any subagent dispatch: confirm the preamble (sub-rule B) is pasted verbatim.

- Before any emulator/device action: confirm the run routes through the Containers submodule CLI, not host-direct; confirm no live device is targeted; confirm the gate host is eligible (otherwise honestly BLOCKED).
- Before any distribute action: confirm Challenge Tests EXECUTED (not compiled) against the exact artifact; version code bumped; CHANGELOG entry present; debug-stage evidence present if two-stage distribute applies.
- Before any push / destructive git action: confirm no force flags without per-operation operator approval; remote is the approved set only; for history rewrite, hardlinked `.git` backup made first.
- Before any host-affecting command: confirm the command is not in the host-power blocked class.

Anti-bluff requirement: the checklist MUST be present in the preamble document and verifiable by `check-constitution.sh` . Each checklist item is traceable to a governing clause cited inline.

Mechanical enforcement.

Consuming-project `check-constitution.sh` (or `verify-all-constitution-rules.sh`) MUST verify ALL of:

1. `constitution/scripts/hooks/guard-forbidden-commands.sh` exists at the canonical path in the constitution submodule.
2. The consumer's AI-agent settings file (`.claude/settings.json` or equivalent) contains a `PreToolUse` hook entry referencing the canonical script path.
3. `docs/AGENT_GUARDRAILS.md` (or the equivalent preamble path) exists, contains the `SUBAGENT CONSTITUTIONAL PREAMBLE` heading, and contains the `ORCHESTRATOR PRE-ACTION CHECKLIST` heading.
4. A hermetic test for the hook exists (at minimum, a file under `tests/hooks/` or equivalent) and passes.
5. The anchor literal `11.4.109` is present in every per-scope governance file (`CLAUDE.md` , `AGENTS.md` , `QWEN.md`) of the consuming project and its owned submodules per §11.4.35 inheritance.

Pre-build gate `CM-ANTI-FORGETTING-ENFORCEMENT` enforces (1)–(4). Propagation gate `CM-COVENANT-114-109-PROPAGATION` enforces (5). Paired §1.1 meta-test mutations: (a) remove the `PreToolUse` hook entry from the agent settings file → gate (2) FAILs; (b) delete `docs/AGENT_GUARDRAILS.md` → gate (3) FAILs; (c)

remove the canonical hook script from the constitution submodule → gate (1)
FAILs; (d) strip the 11.4.109 literal from a consumer CLAUDE.md → propagation
gate FAILs.

Inheritance.

Applies recursively to every submodule, every feature, every new artifact, every project consuming this constitution. Submodule constitutions MAY add stricter rules (additional blocked command classes, additional checklist items, stricter test counts) but MUST NOT relax any clause. Per §11.4.35: this universal clause lives in the constitution submodule; consuming projects' CLAUDE.md / AGENTS.md / QWEN.md carry the inheritance pointer-block per §11.4.35 + the 11.4.109 literal per the propagation gate.

The principle this clause enshrines (in the fewest possible words): A constraint that depends on an agent remembering it is not a constraint — it is a hope. The PreToolUse hook converts the most dangerous command classes into hard-blocked failures regardless of agent memory. The preamble document converts the full ruleset into something the orchestrator can inject rather than recite from training. Together they are the minimum viable "anti-forgetting" enforcement posture for any AI-agent-assisted project under this constitution.

Composes with §11.4 + §11.4.1..§11.4.16 (anti-bluff covenant — the hook and preamble are the floor and ceiling of the anti-bluff enforcement layer), §11.4.6 (no-guessing — the escape hatch requires a <reason> exactly because undocumented exceptions are a guessing posture), §11.4.10 (credentials — the host-power class is categorically non-overridable, analogous to credential leaks), §11.4.75 (mechanical enforcement — §11.4.109 is the AI-agent-side specialisation of §11.4.75's git-hook mechanical enforcement; together they cover the full automated-and-agentic surface), §11.4.76 (Containers mandate — the emulator gate in the hook enforces §11.4.76 at the tool-call boundary), §11.4.78 / §11.4.79 / §11.4.80 (CodeGraph — the hook passes codegraph MCP calls through untouched; the preamble includes the anti-forgetting constraint about CodeGraph use), §11.4.81 (cross-platform parity — the hook script is portable bash/awk with no OS-specific primitives; per-OS blocked-class equivalents are addable via consuming-project extension), §11.4.84 (working-tree quiescence — the pre-action checklist instructs agents to grep for mutation markers before staging), §11.4.98 (full-automation mandate — the hook's test suite MUST itself be fully self-driving end-to-end), §11.4.102 (systematic-debugging — any hook violation that surfaces

should trigger the systematic-debugging arc, not a quick workaround), §12 (host-session-safety — the host-power class is non-overridable; §12 is the constitutional root of that prohibition).

Classification: universal (§11.4.17) — the PreToolUse guard hook + subagent preamble + orchestrator checklist pattern is vendor-neutral (defined against the JSON stdin contract that every Claude Code hook receives; the pattern ports to any agent runtime that exposes a pre-tool-call interception point) and reusable across ANY project that uses AI coding agents. The specific blocked command-class list at the UNIVERSAL layer covers the four classes that are unconditionally forbidden by this constitution (host-power, privilege escalation, force-push/bypass, raw emulator host-direct); consuming projects extend the list with project-specific additions per §11.4.35.

Non-compliance is a release blocker. No escape hatch — no `--skip-pretooluse-hook`, `--no-guardrails-doc`, `--anti-forgetting-optional`, `--single-layer-sufficient` flag exists. A project whose agent can forget a constitutional constraint and execute the forbidden command is constitutionally undefended, irrespective of how well the docs are written.

Canonical authority: this Constitution.md §11.4.109 in the HelixConstitution submodule; reference implementation `constitution/scripts/hooks/guard-forbidden-commands.sh` + reference preamble document `constitution/docs/AGENT_GUARDRAILS.md` (path within the submodule). Consuming projects inherit by reference per §11.4.80.

2026 06 03

§ . . — Pre-build build-readiness
verdict + change-impact clash detection mandate
(operator mandate, — —)

Short tag: `pre-build-readiness-verdict` .

§-slot history note. Drafted as §11.4.109 in `docs/research/prebuild_rigor_20260603/PROPOSED_CONSTITUTION_ANCHOR.md`, renumbered to §11.4.110 per §11.4.71 fetch-before-push after the concurrent landing of §11.4.109 (Anti-Forgetting Enforcement, commit `1d9e5d6`). Same collision-resolution pattern as §11.4.75 / §11.4.76.

Forensic anchor (genericised, 2026-06-03). A fix shipped a new system-property *read* in a code/init change but introduced no matching security-policy grant and no property-context type entry for it — so the read was *silently denied* at runtime and the feature was dead, while every pre-build gate stayed green because the gate only grepped the source file for the change. The defect class generalises: a change introduces a new dependency on a *second* artifact (security policy, context file, service registry, interface freeze-snapshot, symbol table, build-graph node) that the pre-build never cross-checks, so the change is "ready" by every existing gate yet broken the instant it runs. This is the SOURCE → ARTIFACT half of §11.4.108's four-layer gap, shifted *left* to pre-build time: most of these clashes are **statically catchable from the change diff itself, before any build**.

The architectural flaw (FACT, traced). Pre-build gates that validate each artifact *in isolation* (the policy file is internally consistent; the freeze-snapshot is well-formed) do NOT cross-check that a change introducing a new dependency *also* introduced the matching second artifact. The cross-check is **diff-driven** (its input is the change set, not any single file) and is therefore a capability no per-file gate provides. Without it, the only protection is author-vigilance ("remember to also add the policy line"), which §11.4.75 already established is the failure mode that mechanical enforcement exists to replace.

The mandate (ALL must hold).

1. **A single deterministic READY-FOR-BUILD verdict is mandatory.** Before any artifact rebuild, the project MUST emit ONE machine-checkable verdict — *the codebase IS / IS NOT ready for build* — backed by a captured-evidence bundle enumerating every gate family's status. The rebuild orchestrator MUST refuse to start the build unless the verdict is READY (non-zero exit on NOT-READY). A green individual gate is necessary but not the verdict; the verdict is the aggregate function and cannot report READY while any family reports FAIL.
2. **A diff-driven change-impact + conflict/clash detector is mandatory.** For every change in the batch, the pre-build MUST cross-check that any newly-introduced dependency on a *second* artifact is satisfied — at minimum: a new system-property read ⇔ the property is typed in the platform's property-context registry AND the reading domain has the matching read-grant; a new registered service ⇔ its service-context entry exists; a new declared startup/init service ⇔ its security label is present; a new/changed stable interface ⇔ its freeze-snapshot/version was updated in the same change; a new native-library

dependency $\Rightarrow$ the named library resolves to a known module or prebuilt; a new security-policy rule $\Rightarrow$ every type/attribute it uses is defined; and two changes in the batch touching the SAME function / resource / configuration seam $\Rightarrow$ the collision is explicitly acknowledged (the cross-fix-interaction class). Each sub-detector is independently anti-bluff (clause 5).

3. **Coverage-completeness is a gate, not a convention.** Every changed source file in the batch MUST map to ≥ 1 pre-build gate + ≥ 1 deployed-target test + ≥ 1 paired §1.1 mutation (the four-layer §11.4.4(b) floor). A baseline file records the legacy tail of known holes; the gate FAILs on any NEW hole and the covered-ratio ratchets upward over phases per §11.4.50. A change with no owning gate/test/mutation is a coverage hole and blocks the verdict.
4. **Two-speed honesty — grep-speed vs REQUIRES_BUILD — is mandatory.** The regime MUST separate always-on grep-speed gates (run every commit, seconds) from REQUIRES_BUILD heavy gates (the build-graph parse-only dry-run; full security-policy neverallow compilation; built-symbol ABI diff). Heavy gates run as diff-gated opt-in stages (triggered only when the change touches build-graph / policy / interface files) and MUST run bounded per the host-session-safety budget (§12.6/§12.7) or inside the project's containerized build. Conflating the two — making the grep path slow, or claiming grep proves graph/neverallow correctness — is forbidden per §11.4.6.
5. **Every gate and analyzer is anti-bluff by paired mutation.** Each gate family AND each wired third-party analyzer (its exit-code-on-failure semantics included) MUST ship a paired §1.1 mutation that plants the exact defect class and asserts the gate FAILs (e.g. a property read with no policy $\rightarrow$ the clash detector FAILs; a syntactically-broken build-graph fragment $\rightarrow$ the dry-run reports FAIL; an analyzer that exits 0 on a planted violation is itself the bluff the mutation catches). An analyzer wired without its mutation pair is a §11.4 bluff surface.
6. **Honest boundary — pre-build catches statically-catchable defects only.** A READY verdict proves the change is internally consistent and ready to build; it does NOT prove the feature works on the deployed target. The correct claim is: *after the pre-build passes, the residual defects deployed-target / manual testing can find are runtime/hardware defects, NOT statically-catchable (preventable) ones* — the regime makes the **preventable** class empty, not the **all-defects** class. Overclaiming that manual testing will find nothing is itself a §11.4.6

violation. Build-graph correctness, full neverallow safety, and ABI correctness are proven ONLY by the REQUIRES_BUILD families (clause 4); the grep tier proves harness-presence + symbol-definedness, never full correctness.

Runtime/USER-VISIBLE correctness remains entirely §11.4.108's four-layer job.

Classification: universal (§11.4.17) — the single ready-for-build verdict, the diff-driven change-impact/clash detector, the coverage-completeness gate, the grep-speed-vs-REQUIRES_BUILD two-speed split, the per-gate paired mutation, and the statically-catchable-only honest boundary are platform-neutral pre-build-rigor disciplines reusable by ANY project that builds an artifact from source; the consuming project supplies its concrete property/service/policy registries, its build-graph parse-only command, its interface-freeze mechanism, and its changed-file-to-gate mapping per §11.4.35.

Composes with §11.4.1 (FAIL-bluffs forbidden — a clash-detector that crashes on a malformed diff and FAILs is as misleading as one that passes silently), §11.4.4 (test-interrupt-on-discovery — a NOT-READY verdict triggers STOP-and-fix before the build), §11.4.6 (no-guessing — the two-speed honest boundary; the grep tier MUST NOT claim correctness it cannot prove), §11.4.9 (batch-source-fixes-before-rebuild — the verdict gates exactly that rebuild), §11.4.27 (no-fakes-beyond-unit — wired analyzers exercise the real change), §11.4.50 (deterministic consistency — the coverage-completeness ratchet shares §11.4.50's threshold mechanism and N-iteration determinism), §11.4.67 (target-shell-parseability — the clash-detector helper itself parses under its target shell), §11.4.75 (mechanical enforcement — §11.4.110 is the mechanical replacement for author-vigilance on the clash class), §11.4.92 (multi-pass change-evaluation — the verdict is the mechanical companion to Pass 1/2/3), §11.4.108 (four-layer fix-verification — §11.4.110 is the SOURCE → ARTIFACT half shifted left to pre-build time; §11.4.108 owns the RUNTIME-ON-CLEAN-TARGET → USER-VISIBLE half; together they span all four layers).

Propagation. Propagation gate `CM-COVENANT-114-110-PROPAGATION` enforces the literal anchor `11.4.110` across the consumer fleet (parent + owned-submodule `CLAUDE.md` / `AGENTS.md` / `QWEN.md`); paired §1.1 meta-test mutation strips the `11.4.110` literal from a consumer file → the gate FAILs. Recommended per-family gates `CM-READY-FOR-BUILD-VERDICT` (single deterministic verdict gates the rebuild), `CM-CHANGE-IMPACT-CLASH-DETECTOR` (diff-driven second-artifact cross-check), `CM-COVERAGE-COMPLETENESS-GATE` (changed-

file $\Rightarrow$ gate+test+mutation with baseline ratchet), CM-BUILDGRAPH-DRYRUN-WIRED + CM-SEPOLICY-NEVERALLOW-WIRED (REQUIRES_BUILD opt-in stages exist + bounded-safe); each with a paired §1.1 mutation per clause 5. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md §11.4.110 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no

```
--source-green-is-ready , --skip-clash-detector , --skip-coverage-  
gate , --no-ready-verdict , --grep-proves-neverallow , --skip-  
buildgraph-dryrun , --build-without-ready-verdict flag exists.  
2026 06 03
```

§ . . — **Resolve-by-stable-name-not-by-enumeration-index** mandate (research-derived,
- -)

Short tag: resolve-by-stable-name .

Forensic anchor (genericised, 2026-06-03). A platform bound an audio output to a kernel-enumerated device index (card=0). A second device of a different class enumerated FIRST at boot and took slot 0, shifting the intended device to slot 1. The static index binding now pointed at the wrong device: the policy layer failed to attach the intended sink, mis-assigned its TYPE, and the user-facing output switcher labelled the AV receiver "Wired Headphone" while routing collapsed to stereo. The lower layer of the SAME stack already resolved the device correctly **by name** (scanning /proc/asound/card*/id for the controller name) — proving the brittleness was the *index* binding, not the resolution capability. The defect class generalises to every enumerated resource whose ordinal is assigned at discovery/ boot/hotplug time and is therefore non-deterministic across reboots, device additions, and topology changes.

The mandate. Any binding between a configuration / policy / code site and a hardware device, resource handle, or enumerated entity — audio cards, display connectors, network interfaces, storage devices, GPU render nodes, input devices, camera indices, container/process slots, any registry whose members are assigned an ordinal at discovery time — **MUST** resolve the target by a **stable identifier** (name, UUID, serial, label, controller-name, content-derived hash, sink-reported identity) and **MUST NOT** resolve it by an enumeration index / ordinal / slot number,

UNLESS the platform documents that ordinal as deterministically pinned (and the pinning mechanism is itself part of the binding, captured and asserted). Where a stable identifier exists at one layer of a stack, every other layer binding the same resource MUST use the same identifier — mixed by-name-here / by-index-there is the structural weak link this anchor forbids. Where the platform offers a sink-reported / capability-derived identity (e.g. an EDID/ELD-style capability descriptor), preferring it makes the binding robust to the broadest class of topology change.

Honest boundary. When the platform genuinely provides no stable identifier and only an ordinal exists, the binding MUST (a) pin the ordinal deterministically via the platform's own pinning mechanism, (b) capture that pin as part of the binding's runtime signature per §11.4.108, and (c) document the residual fragility as

UNCONFIRMED: -class risk per §11.4.6 — never silently trust an unpinned ordinal.

Classification: universal (§11.4.17) — index-vs-name resource resolution is a platform-neutral binding-robustness discipline reusable by ANY project that binds to enumerated hardware / resources / handles; the consuming project supplies its concrete stable-identifier mechanism (ALSA card name, DRM connector name, NIC predictable name, block-device UUID, etc.) and the specific layers that must agree per §11.4.35.

Composes with §11.4.6 (no-guessing — "the index is *usually* stable" is the exact guess this anchor forbids; prove the pin or bind by name), §11.4.8 (deep-web-research — mature platform stacks resolve by name/UUID; reproducing a known-brittle index binding when the by-name path is documented is a §11.4.8 violation by omission), §11.4.69 (sink-side positive-evidence — the by-name binding's correctness is verified by the downstream/sink-reported identity, not by the local config), §11.4.108 (four-layer fix-verification — the resolved-by-name binding's runtime signature is asserted on a clean target across the topology that originally broke the index binding), §11.4.110 (pre-build clash detection — a new ordinal-based binding in a diff is a statically-catchable clash class the change-impact detector flags).

Propagation. Propagation gate CM-COVENANT-114-111-PROPAGATION enforces the literal anchor 11.4.111 across the consumer fleet; paired §1.1 meta-test mutation strips the literal → the gate FAILs. Recommended per-family gate CM-RESOLVE-BY-NAME-NOT-INDEX (diff-driven scan flagging a new enumeration-index binding to a resource that exposes a stable identifier, unless an accompanying

deterministic-pin + runtime-signature is present); paired §1.1 mutation plants an unpinned index binding → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md §11.4.111 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no

```
11 4 112
--allow-index-binding , --ordinal-is-stable-enough , --skip-name-
resolution , --trust-unpinned-index flag exists.
2026 06 03
```

§ . . — **Structural-impossibility won't-fix classification mandate (research-derived, - -)**

Short tag: `structural-impossibility-wont-fix` .

Forensic anchor (genericised, 2026-06-03). Deep research (§11.4.8) into relocating protected (content-protection / secure-surface) video to a secondary display proved — from authoritative platform/HDCP documentation plus reproducible captured behaviour (a secure surface is blanked on any output lacking the secure flag; mirror/screenshot of a secure layer returns black) — that the goal is **structurally impossible by platform design**, not a missing feature or an unsolved engineering problem. Without a durable classification, such a goal is re-investigated every cycle: each new agent / operator re-reads the same sources, re-runs the same probes, re-derives the same impossibility, and burns the same effort — an anti-pattern that compounds across cycles.

The mandate. When deep research per §11.4.8 PROVES (with cited authoritative sources AND, where applicable, reproducible captured evidence) that a goal is structurally impossible on the target platform — i.e. forbidden by the platform's design / a hardware/protocol constraint / a documented kernel-or-API limitation, not merely unimplemented or hard — the goal MUST be: (1) **classified won't-fix and closed** per the §11.4.90 terminal-status discipline, with obsolescence/closure reason `structurally-impossible` (a value in the §11.4.90 closed reason vocabulary); (2) **documented with the impossibility evidence** — the cited authoritative source URLs (per §11.4.99 latest-source verification where the platform docs may evolve) PLUS the reproducible probe/captured evidence that demonstrates the constraint, recorded in the tracker entry and the relevant `docs/`

guide so the verdict is auditable; (3) **NOT re-attempted in future cycles** — the closed entry is the canonical answer; a future agent that re-opens it **MUST** cite **NEW** evidence that the platform constraint changed (per §11.4.34 reopened-source attribution + §11.4.7 demotion-evidence), never merely re-derive the same impossibility; (4) **paired with the correct posture** — the entry **MUST** state what the project **DOES** instead (the supported-within-the-constraint behaviour), so "impossible" is never confused with "broken / unhandled".

Honest boundary (§11.4.6). `won't-fix: structurally-impossible` is reserved for *proven* platform/hardware/protocol impossibility. "We could not find a way" / "it is very hard" / "no time" are NOT structural impossibility — they are `operator-blocked` (§11.4.21) or open work, and mislabelling them `won't-fix` to avoid the work is a §11.4 bluff at the planning layer. A future platform change **CAN** make the impossible possible; the classification is durable but not eternal, and reopening on **NEW** cited evidence is correct.

Classification: universal (§11.4.17) — proving-then-durably-classifying structural impossibility is a platform-neutral effort-conservation + honesty discipline reusable by ANY project; the consuming project supplies the specific impossible goal, its platform constraint, and the cited evidence per §11.4.35.

Composes with §11.4.6 (no-guessing — the impossibility is stated as **FACT** with cited evidence, never as "probably can't"), §11.4.7 (demotion-evidence — reopening a `won't-fix` requires **NEW** positive evidence the constraint changed), §11.4.8 (deep-web-research — the proof is the research output; "NO external solution found — structurally impossible" is the §11.4.8 citation), §11.4.34 (reopened-source attribution — a reopen of a structurally-impossible item attributes who reopened + the new evidence), §11.4.90 (Obsolete/terminal-status — `structurally-impossible` is a closure reason in the §11.4.90 closed vocabulary; the colorizer marks it like any closed item), §11.4.99 (latest-source verification — the cited platform docs **MUST** be the latest where they can evolve, so the impossibility verdict is not stale).

Propagation. Propagation gate `CM-COVENANT-114-112-PROPAGATION` enforces the literal anchor `11.4.112` across the consumer fleet; paired §1.1 meta-test mutation strips the literal → the gate **FAILs**. Recommended per-family gate `CM-WONT-FIX-STRUCTURAL-IMPOSSIBILITY` (every tracker entry with closure reason `structurally-impossible` carries cited authoritative-source evidence + a stated

correct-posture line); paired §1.1 mutation strips the evidence citation from such an entry → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

(5) SCOPE-STATEMENT + ADJACENT-GOAL ENUMERATION — an impossibility verdict is bounded, and MUST say where its edge is (extension, research-derived, 2026-07-17). Forensic FACT (genericised, from a consuming project): a verdict that was TRUE on its own terms — "relocating another application's protected video output to a secondary display **while that application's UI stays in place, without the application's cooperation** is structurally impossible" — LEAKED to an ADJACENT, VIABLE goal it never covered: moving the whole task to the other display. For six weeks the project believed "no platform mechanism exists", while the platform had shipped a public, permission-free SDK call for the adjacent goal several major versions earlier, and the platform vendor's own connected-display documentation PRESCRIBED exactly that call. The impossibility was never re-litigated (clause (3) worked); it was CITED, unexamined, against a neighbour — and clauses (1)-(4) contain nothing that bounds a verdict's reach. This anchor is ASYMMETRIC without this clause: clause (3) hardens a verdict against reopening ("the closed entry is the canonical answer") while nothing constrains how far it may be applied, so the only mechanical pressure runs in the direction that AMPLIFIES a scope leak. Therefore every structurally-impossible verdict MUST additionally: **(a) STATE ITS EXACT SCOPE** — the precise goal/variant the cited evidence actually proves impossible, written as the narrowest claim the evidence supports (every qualifier that is load-bearing — "without the application's cooperation", "while the UI stays put", "on this platform version" — is part of the verdict, not commentary; a qualifier dropped is a verdict widened); **(b) ENUMERATE THE ADJACENT GOALS IT DOES NOT COVER** — the neighbouring variants a reader could mistake for the same goal, each marked NOT-COVERED-BY-THIS-VERDICT (a neighbour whose viability is unknown is recorded as UNKNOWN with a tracked follow-up per §11.4.6 — an unexamined neighbour is NEVER silently absorbed into the impossibility); **(c) BE CITATION-FENCED** — the verdict MUST NOT be cited as the answer for ANY goal outside its stated scope, in planning, investigation, design, or status, without its OWN §11.4.150 deep-research pass and its OWN proof (citing it beyond its fence is asserting an unproven claim as fact — a §11.4.6 violation and a §11.4/§11.4.1 bluff at the verdict layer, of the same class §11.4.199 names when a reproduction that never reaches the precondition is used to conclude "the mechanism never

engages"). **Honest boundary (§11.4.6):** an adjacent goal being outside the fence does NOT make it viable — it makes it UNDECIDED, and it earns its own verdict by its own evidence, never by inheritance in either direction. **Why this is a clause and not another narrative:** §11.4.150(A)/(E) already documented this exact failure class in prose (its case study is an impossibility verdict OVERTURNED by the research pass it never had) and REQUIRES the deep-research pass before a `structurally-impossible` verdict is EARNED — yet the leak still happened, because §11.4.150 fires at the seam where a verdict is MINTED for a tracked item, and this leak happened where a verdict was CITED as a premise, which crosses no closure seam at all. Clauses (2) and (4) are already content requirements on the verdict artifact; (5) is the third, and it is the one that bounds blast radius.

Gate contract for clause (5) (extension, 2026-07-17).

`CM-WONT-FIX-STRUCTURAL-IMPOSSIBILITY` is EXTENDED: every `structurally-impossible` entry MUST additionally carry (i) an explicit SCOPE statement and (ii) an adjacent-goals-NOT-covered enumeration (each neighbour marked not-covered or UNKNOWN + tracked); and (iii) a citation of the verdict against a goal outside its stated scope, without that goal's own §11.4.150 research + proof, → FAIL. Paired §1.1 mutations: strip the scope statement, or the adjacent-goal enumeration, from an impossibility entry → the gate FAILs; cite an in-scope verdict against an enumerated not-covered neighbour → the gate FAILs. Gate-code = a separate work item; this is the contract, not a claim the code has shipped (§11.4.6).

Canonical authority: this Constitution.md §11.4.112 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no

`--wont-fix-without-proof` , `--reattempt-closed-impossible` , `--skip-impossibility-evidence` , `--impossible-equals-broken` , `--unscoped-impossibility-verdict` , `--skip-adjacent-goal-enumeration` , `--cite-verdict-beyond-scope` , `--neighbour-inherits-impossibility` flag exists.

§ . . — Absolute no-force-push +
merge-onto-latest-main mandate (User mandate,
- -)

Short tag: `absolute-no-force-push` .

Forensic anchor — verbatim user mandate (2026-06-03):

"Any force-push is strictly forbidden! We must for every Submodule take as a base latest commit on Submodule's main (or master) branch, then on top of it carefully to merge all changes that have to be pushed! Once all merging is carefully done we perform commit and push to all Submodule's upstreams!"

The mandate — force-push is STRICTLY FORBIDDEN, with NO exception. Any force-push — `git push --force` , `git push --force-with-lease` , `git push +<ref>` , or any history-rewriting operation that overwrites a remote ref non-fast-forward — is forbidden against every repository this Constitution governs (every consuming project's main repo, this constitution submodule, every owned submodule, every nested submodule, every upstream of each). There is no escape hatch, no operator-approval path, no "after a merge-first audit" path. The merge-onto-latest-main procedure below is ALWAYS available, so a force-push is NEVER necessary; eliminating the need eliminates the operation.

The mandated 6-step integration procedure (for every repo/submodule whose local has commits to publish OR whose mirrors have diverged):

1. **Fetch all remotes** — `git fetch --all --prune --tags` against `origin` + every configured upstream; capture the output as evidence.
2. **Set the base to the LATEST commit on the canonical `main / master` branch** — the most-advanced mirror tip across all upstreams. The integration target is that tip, never a stale local branch point.
3. **Carefully MERGE every change that must be published on top of that base** — a union merge that preserves BOTH sides (local commits AND every remote-side commit). NEVER `-s ours` (it silently discards the remote side), NEVER a rebase/reset/ `--allow-unrelated-histories` that could drop commits (per §9 no-commit-loss). The merge commit MUST descend from every mirror tip so that every subsequent push is a clean fast-forward.

4. **Resolve every conflict carefully** so nothing is broken, lost, or corrupted on either side — no conflict markers anywhere (`grep -rn '^<<<<<< \|^===== $\|^>>>>>> '` returns empty across governance + source + test files), no file silently dropped, every previously-passing gate/test still passes on the integrated tree (per §11.4.4 + §11.4.40), every captured-evidence artefact still validates.
5. **Commit the merge** — stage only the intended files (NEVER `git add -A` inside a submodule with untracked build content per §11.4.30), commit the carefully-resolved merge.
6. **Push the result to ALL upstreams** — each push is a fast-forward because the merge commit descends from every mirror tip, so NO force is ever needed. If an upstream still reports non-fast-forward (a sibling landed work between step 1 and the push), do NOT force — return to step 1 for that upstream, fetch its new tip, MERGE it (union, preserve both), re-validate step 4, and re-push.

Tightens §11.4.41 / §11.4.71 / §9.2 / CONST-043 — force-push escape hatch

REMOVED. §11.4.41 (Pre-Force-Push Merge-First Mandate), §11.4.71 (Pre-Push Fetch + Investigate + Integrate), §9.2 (Force-push requires explicit user authorization), and CONST-043 previously PERMITTED a force-push after a merge-first pipeline + per-operation operator approval. §11.4.113 supersedes that permitting stance: even WITH operator approval, even AFTER a clean merge-first audit, a force-push is forbidden — because the merge-onto-latest-main path in this clause is always available and always sufficient, force is never necessary and therefore never authorised. Those clauses' merge-first/fetch-first machinery REMAINS in force as the integration discipline; only their terminal "...then force-push" step is struck. §9.2's destructive-operation §9 backup discipline still applies to any non-push history operation, but a history-rewriting *push* is simply not performed.

No escape hatch. No `--force`, no `--force-with-lease`, no `+<ref> push`, no `--no-verify` to slip a forced ref past a hook, no "operator authorised the force-push" path, no "lease semantics unavailable so plain --force is OK" path. A force-push that lands on any remote is a critical defect regardless of how briefly it persisted, severity-equivalent to a §11.4 PASS-bluff at the data-safety layer (per §9).

Classification: universal (§11.4.17) — an absolute-no-force-push + merge-onto-latest-main integration discipline is platform-neutral and reusable across every repository in every consuming project; nothing here references a particular project, vendor, hardware, or region.

Composes with §2.1 (multi-upstream push is the norm — step 6 fans out to all), §9 / §9.2 (absolute data safety — §11.4.113 is the no-loss push discipline that makes force-push unnecessary), §11.4.4 (test-interrupt + clean-baseline retest — step 4 audit), §11.4.6 (no-guessing — remote state is not knowable without the step-1 fetch), §11.4.26 (constitution-submodule update workflow — its conflict-resolution step is this procedure), §11.4.37 (fetch-before-edit — step 1 generalised to fetch-before-push), §11.4.40 (full-suite retest authority for the integrated tree), §11.4.41 (force-push merge-first — TIGHTENED: merge-first stays, the force-push step is removed), §11.4.71 (pre-push fetch + integrate — TIGHTENED the same way), §11.4.88 (background-push — the detached push still follows this fast-forward-only discipline), CONST-043 (per-operation authorization — TIGHTENED: no force-push is authorisable).

Propagation. Propagation gate `CM-COVENANT-114-113-PROPAGATION` enforces the literal anchor `11.4.113` across the consumer fleet; paired §1.1 meta-test mutation strips the literal → the gate FAILs. Recommended per-family gate `CM-NO-FORCE-PUSH-ABSOLUTE` (scans tracked scripts + hooks for any `push --force / push --force-with-lease / push +<ref>` invocation and rejects it; a §11.4.109-class `PreToolUse` guard blocks the force-push command class at the tool-call boundary); paired §1.1 mutation injects a `git push --force` into a tracked script → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md §11.4.113 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--force` , `--force-with-lease` , `--allow-force-push` , `--force-push-authorised` , `--skip-merge-onto-main` flag exists.

§ . . — Last-known-good-tag
 regression isolation mandate (. . -dev
 remediation, - -)

Short tag: last-known-good-tag-isolation .

Forensic anchor (genericised, 2026-06-03). A user reported that a previously-working feature (multichannel HDMI audio) was broken on the current build. The operator's lead — "the feature WORKS and passes all tests at the last release tag" — converted an open-ended root-cause hunt into a targeted regression repair: diffing the broken working tree against the last-known-good tag's commit isolated EVERY reported audio defect (multichannel downmix + output-device label mismapping) to ONE post-tag burndown batch
 (`AudioDeviceInventory.connectDualHdmiPorts_l` dual-HDMI-connect) in minutes, not hours. The tagged-good version is the oracle: behaviour that works there and is broken now is, by definition, a regression introduced in the diff between the two.

The mandate. When a previously-working feature, behaviour, or capability is observed broken, the FIRST diagnostic action MUST be to identify the last release tag (or last commit / build / deployment) at which that feature was KNOWN-GOOD, and diff/bisect the broken state against it — BEFORE any open-ended root-cause investigation, BEFORE any speculative fix. The known-good revision is the regression oracle: (1) it bounds the search space to the commits between known-good and now (a `git diff <good-tag>..HEAD --stat` of the feature's files is the captured evidence); (2) it tells you the fix is a regression REPAIR (restore/forward-fix the broken sub-part), not a from-scratch design problem; (3) each suspect file is compared against its known-good version as the behavioural oracle. When the operator volunteers a known-good tag, that lead is load-bearing and MUST be acted on first. Where the bisect surface is large, `git bisect` against the feature's autonomous RED test (per §11.4.115) mechanises the search.

Forward-fix vs wholesale-revert. Once the regressing change is isolated, the default is a SURGICAL forward-fix — keep the new features the post-good-tag work introduced, revert ONLY the broken sub-part where the known-good version was correct — UNLESS the operator prefers wholesale revert. A wholesale revert of a multi-feature batch to reach the known-good state is a last resort that loses the

batch's other (working) features; prefer isolating and repairing the single regressing seam. Every forward-fix MUST prove (via the §11.4.115 RED → GREEN test) that the feature now matches the known-good behaviour.

Honest boundary (§11.4.6). "It worked before" is a HYPOTHESIS until the known-good tag is identified AND the feature is confirmed working there (re-run the feature's test against the known-good build, or cite the captured-evidence from when that tag shipped). A feature that NEVER worked is not a regression and this anchor's oracle does not apply — investigate as new work. "Probably regressed in the last batch" without the diff is a §11.4.6 guess.

Classification: universal (§11.4.17) — diff-against-last-known-good is a platform-neutral regression-isolation discipline reusable by ANY versioned project with release tags / commit history; the consuming project supplies its tag-naming convention and per-feature known-good oracle per §11.4.35.

Composes with §11.4.4 (test-interrupt-on-discovery — the isolation runs as the first step of the systematic-debugging Phase 1 the STOP triggers), §11.4.6 (no-guessing — the diff is the FACT-grade evidence that replaces "probably the last batch"), §11.4.7 (demotion-evidence — confirming known-good requires positive evidence at that tag), §11.4.40 (full-suite retest authority — the known-good tag is the baseline a release-gate retest restores toward), §11.4.43 (TDD-fix — the §11.4.115 RED test reproduces the regression the diff isolated), §11.4.102 (systematic-debugging — last-known-good isolation IS the Phase 1 root-cause technique for regressions specifically), §11.4.108 (four-layer fix-verification — the forward-fix's runtime signature is verified against the known-good behaviour on a clean target).

Propagation. Propagation gate `CM-COVENANT-114-114-PROPAGATION` enforces the literal anchor `11.4.114` across the consumer fleet; paired §1.1 meta-test mutation strips the literal → the gate FAILs. Recommended per-family gate `CM-REGRESSION-ISOLATED-AGAINST-KNOWN-GOOD` (every tracker entry classified as a regression cites the last-known-good tag + the diff range it was isolated to); paired §1.1 mutation strips the known-good citation from a regression entry → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md §11.4.114 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--skip-known-good-diff` , `--root-cause-from-scratch` , `--assume-regression-source` , `--wholesale-revert-without-isolation` flag exists.

1 1 8

2026 06 03

§ . . — **RED-baseline-on-the-broken-artifact + polarity-switch mandate** (. . -dev remediation, - -)

Short tag: `red-baseline-polarity-switch` .

Forensic anchor (genericised, 2026-06-03). Per §11.4.43 a fix opens with a RED test. The 1.1.8-dev remediation sharpened HOW: the autonomous RED test was authored to REPRODUCE each live defect on the CURRENT pre-fix build — proving the defect is real, live, and captured (video flips to primary on reconfigure 3/3; audio 5.1 → Arvus PCM2/0 ×3; subtitle absent on secondary) — with a single environment flag (`RED_MODE=1` default) that, set to `0` post-fix, flips the SAME test to GREEN-guard polarity. One test source, two roles: the reproduction-of-the-defect (RED) and the permanent regression-guard (GREEN). No separate happy-path test is authored; the test that catches the bug IS the test that guards against its return.

The mandate. Every §11.4.43 RED test MUST be authored to reproduce the defect on the CURRENT, pre-fix artifact (the actual broken build / deployment), capturing positive evidence per §11.4.5 / §11.4.69 / §11.4.107 that the defect is genuinely present — never a synthetic / hypothetical failure the fix is then written to agree with. The SAME test source MUST carry a single polarity switch (an env flag / parameter, canonical `RED_MODE` , default `1` = reproduce-and-assert-defect-present) that, flipped to `0` post-fix, converts the test into the GREEN regression-guard asserting the defect is ABSENT. The RED run (against the broken artifact) and the GREEN run (against the fixed artifact, per §11.4.108 on a clean target) MUST both be captured as evidence: RED-then-GREEN is the proof the test catches the bug AND the fix closes it. A separate happy-path test that never failed on the broken artifact is forbidden as the primary guard — it demonstrates only that the test agrees with the fix (the §11.4.43 PASS-bluff), not that it catches the defect.

Why one test, two roles. A test authored only as a post-fix happy-path can pass on a build where the feature is silently broken (it asserts the fix's intended state, not the defect's absence). A RED-first test that was OBSERVED failing on the broken

artifact, then flipped to GREEN-guard, is provably sensitive to the exact defect — its RED capture is the falsification record. The polarity switch keeps the two roles in ONE source so they cannot drift apart (a separate happy-path test can be edited to stop testing the real condition; the RED-evidence binds the GREEN guard to the actual defect).

Honest boundary (§11.4.6). If the defect cannot be reproduced on the current artifact (the RED run does not fail), that is a finding, not a license to write a GREEN-only test: either the defect is not present on this build (close per §11.4.7 with the negative evidence) or the test does not exercise the real condition (fix the test). A RED test that "passes" on the known-broken artifact is itself a §11.4.1 FAIL-bluff inverse — it proves the test is blind, not that the feature works.

Classification: universal (§11.4.17) — RED-baseline-on-the-broken-artifact + a single-source polarity switch is a platform-neutral test-construction discipline reusable by ANY project doing TDD-fix; the consuming project supplies its polarity-flag mechanism (env var, build flag, parameter) and per-defect reproduction driver per §11.4.35. Strict refinement of §11.4.43 step RED.

Composes with §11.4.1 (FAIL-bluffs — a RED test that passes on the broken artifact is blind), §11.4.2 / §11.4.5 / §11.4.69 / §11.4.107 (the RED reproduction captures positive defect-present evidence), §11.4.4 (the RED test is the test-interrupt reproduction), §11.4.7 (RED-not-reproducible is a demotion that needs negative evidence), §11.4.43 (refines its RED step — RED is on the real broken artifact, GREEN is the same test flipped), §11.4.50 (deterministic consistency — both RED and GREEN runs are N-iteration stable), §11.4.108 (the GREEN run is the runtime-signature verification on a clean target), §11.4.114 (the RED test mechanises the §11.4.114 bisect against the known-good oracle).

Propagation. Propagation gate `CM-COVENANT-114-115-PROPAGATION` enforces the literal anchor `11.4.115` across the consumer fleet; paired §1.1 meta-test mutation strips the literal → the gate FAILs. Recommended per-family gate `CM-RED-POLARITY-SWITCH-PRESENT` (every regression-fix RED test carries a single polarity switch + a captured RED-on-broken-artifact run before the fix); paired §1.1 mutation removes the polarity switch (splitting RED and GREEN into a happy-path-only guard) → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

(F) Machine-written verdict pairs + guard-viability (extension, research-derived, 2026-07-17). Forensic FACT (genericised, from a consuming project's fix-lifecycle forensics, 2026-07-17): 69 commits in one two-month window moved fix claims/statuses at "GREEN-PENDING-BUILD" — the RED → GREEN discipline existed as prose, but the verdicts were hand-declared, so nothing distinguished a polarity flip that HAPPENED on the target from one that was ASSERTED; the project's own history then selected the discriminator — every fix whose done-claim was preceded by an on-device polarity flip that held recorded ZERO reopens, and every fix confirmed at source-green recurred. Therefore (ALL hold): (1) **RED and GREEN are machine-written verdict artifacts, never prose** — the guard HARNESS (not a human, §11.4.205(4)) writes a per-run verdict file carrying at minimum the item id, guard identity, polarity, exit code, the target's artifact/build fingerprint READ FROM THE TARGET AT RUN TIME, iteration count, and the evidence-file list (path + hash + feature-class); a hand-authored verdict is detectable because the fingerprint/target identity must re-verify against the artifacts it cites. (2) **The pair must cross-validate** — GREEN is acceptable only with a prior RED whose exit $\neq 0$ on the PRE-fix artifact; GREEN's fingerprint $\neq$ RED's fingerprint (an identical fingerprint proves the fix was never deployed); GREEN runs on a clean target per §11.4.108/§11.4.139 with iterations ≥ 3 per §11.4.50. (3) **Guard viability** — a guard that has NEVER been observed FAILing on the genuinely-broken artifact is unvalidated instrumentation (the chaos-engineering validate-the-detector principle): it mints no verdicts; its RED run on the real defect IS its §1.1 mutation evidence (the real defect is the mutant, and the guard provably killed it); corollary: a guard structurally unable to FAIL — or unable to PASS (e.g. a pass-floor above the input population's size) — produces no information in either direction and satisfies nothing. (4) **Mutation = revert + anti-tautology** — the canonical §1.1 mutation for a landed fix is the fix-commit's revert (its semantic inverse, restoring ALL enable-paths of the defect per §11.4.194); a mutation whose diff only deletes/renames the literal strings the gate greps for is a tautology ("grep for X fails after you delete X") and is REFUSED — measured failure: a sed-string mutation "validated" a source-grep gate for a rendered-pixel defect while the defect's second enable-path stayed invisible to both. (5) **Oracle-class** — the verdict's evidence files MUST match the defect's feature-class evidence shape (pixel-class defect $\Rightarrow$ pixel/capture artifacts; audio $\Rightarrow$ sink-side/captured-audio; runtime $\Rightarrow$ runtime observables per §11.4.69); a source-grep transcript can never satisfy a user-visible class — the wrong-layer oracle dies by construction, not by review. Recommended gate `CM-GUARD-VERDICT-MACHINE-WRITTEN` (every done-claiming item's verdict pair is

harness-written, fingerprint-cross-validated, class-matched) + paired §1.1 mutation (forge a verdict whose fingerprint does not re-verify, or pair a gate with a string-deletion mutation → the gate FAILs; gate-code = separate work item). Composes §11.4.6 / §11.4.50 / §11.4.69 / §11.4.107(10) / §11.4.108 / §11.4.135 / §11.4.139 / §11.4.146 / §11.4.194 / §11.4.205.

(G) Defect-identity binding for CONSTRUCTED reproductions — a fix validated only against an investigator-built mechanism is HARDENING, never the fix for the reported defect (extension, research-derived, 2026-07-23).

Forensic FACT (genericised, 2026-07-22): an exit-hang fix (a bounded lock wait + a corrected file-descriptor open) was built and nearly shipped with a confident RED → GREEN pair — but the RED reproduced a mechanism the INVESTIGATOR had constructed (a hand-made process holding the lock), not the operator's reported defect; later systematic debugging (§11.4.102), driving the operator's sequence shape per §11.4.199 with per-step trace timing on a real PTY, REFUTED every mechanism the fix addressed (the lock/unlock pair provably completes on the real path — 6/6 clean interactive exits), and the change was honestly RE-LABELLED in-source as "defensive hardening ... NOT a fix for the reported exit-hang". The gap clauses (A)–(F) leave open: the constructed RED genuinely FAILED on the real pre-fix artifact — this anchor's letter was satisfied — while the MECHANISM it reproduced was the investigator's invention, so the pair proved only that the fix closes the invented mechanism, not the reported defect. Therefore (ALL hold): **(1)** a RED test's PRECONDITION MUST be TRACEABLE to the reported defect's real evidence — the reporter's reproduction sequence (§11.4.199), a captured failure log/trace, or a probe showing the precondition actually occurs in the reported environment — and that traceability is recorded WITH the RED verdict (the clause-(F) fields plus a precondition-provenance field: `observed` vs `constructed`). **(2)** When NO working reproduction of the reported defect exists and the RED's precondition is CONSTRUCTED, the resulting fix mints AT MOST "defensive hardening" — labelled as such in the commit AND the tracker — and NEVER closes the reported defect's item; the item stays OPEN (§11.4.197) until a RED with an `observed` precondition lands, or the item is closed with §11.4.7-grade non-reproduction evidence or a §11.4.90/§11.4.112 evidence-backed reason. **(3)** When systematic debugging REFUTES every mechanism a pending fix addresses, the fix's done-claim for the reported defect is REVOKED on the spot — the change may still land re-labelled as hardening (independently crossing §11.4.142/§11.4.194 review), but claiming it fixes the reported defect after its

mechanisms were refuted is a §11.4/§11.4.1 bluff at the investigation layer even though its RED → GREEN pair is real. STRENGTHENS clause (A)'s "never a synthetic failure" from artifact-synthetic to MECHANISM-synthetic, and §11.4.199 for the no-working-repro case its clause (1) cannot reach (there is no existing sequence to deviate FROM). Recommended gate `CM-RED-PRECONDITION-DEFECT-TRACEABLE` (every defect-CLOSING claim's RED verdict carries precondition-provenance `observed` + its evidence path; `constructed` provenance on a defect-closing claim → FAIL, while the SAME provenance on an item honestly labelled hardening PASSES — the §11.4.201(1) false-positive guard) + paired §1.1 mutation (flip a `constructed` -provenance item from hardening to a defect-closing status → the gate FAILs; gate-code = separate work item). Composes §11.4.6 / §11.4.7 / §11.4.102 / §11.4.142 / §11.4.146 / §11.4.197 / §11.4.199 / §11.4.201.

Canonical authority: this Constitution.md §11.4.115 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--green-only-test`, `--skip-red-reproduction`, `--separate-happy-path-suffices`, `--synthetic-red-OK` flag exists.

44

1 1 8

2026 06 03

§ . . — Real-time
conductor autonomous-test-framework sync
channel mandate (. . -dev remediation,
- -)

Short tag: `conductor-test-sync-channel` .

Forensic anchor (genericised, 2026-06-03). An autonomous test / QA framework (HelixQA) was wired into the conductor's loop. For the conductor (the orchestrating agent / operator) to stay in real-time sync with a long-running autonomous session — knowing which phase / challenge is executing, what evidence was captured, which LLM/vision call ran, and the per-item PASS/FAIL/SKIP/OPERATOR-BLOCKED verdict the moment it lands — the framework exposed a structured append-only JSONL event stream (`conduit.events.jsonl`) plus an atomically-rewritten status snapshot (`conduit.status.json`) that the conductor tails live.

Without it, the conductor either polls blindly (idle waiting, §11.4.94 violation) or learns the outcome only at session end (no mid-session course-correction, no §11.4.4 test-interrupt-on-discovery).

The mandate. Any autonomous, long-running test / QA / validation framework that an external orchestrator (conductor agent or operator) depends on for real-time decisions MUST expose a real-time sync channel with two parts: (1) a **structured append-only event stream** (JSONL or equivalent, one event per line, never rewritten) emitting at minimum session-start / phase-transition / per-test-or-challenge-start / captured-evidence-path / external-call (LLM / vision / sink-probe) / error / per-item-verdict events; (2) an **atomically-rewritten status snapshot** (a single small JSON/file written via write-temp-then-rename so a reader never observes a torn write) carrying the current session/phase/item + running counters + last verdict. Verdicts MUST use the closed vocabulary PASS / FAIL / SKIP / OPERATOR-BLOCKED (per §11.4.45 status vocabulary). The conductor tails the stream / snapshot live (`tail -f` / a dedicated monitor) so it stays in real-time sync, can §11.4.4-interrupt on a fresh defect, and never idles blindly per §11.4.94 / §11.4.97. The channel is the framework's contract with the orchestrator, not an afterthought log.

Anti-bluff for the channel itself (§11.4 / §11.4.69). The event stream + snapshot are themselves captured-evidence artefacts: a verdict event MUST carry the evidence path that backs it (per §11.4.69 `ab_pass_with_evidence`), so a PASS event with no evidence path is a §11.4 PASS-bluff at the channel layer. A status snapshot that reports PASS while the event stream shows no captured-evidence event for that item is a contradiction the conductor MUST treat as FAIL. The channel never invents progress: an item with no start-event cannot have a verdict-event.

Decoupling (§11.4.28). When the framework is an owned, project-agnostic submodule, the channel MUST stay project-neutral — the consuming project registers its own data (endpoints, app/package identifiers, sink hosts) at runtime via the framework's public API, never hardcoded into the framework. The channel format is generic; the payloads are consumer-supplied.

Classification: universal (§11.4.17) — a real-time append-only event stream + atomically-rewritten status snapshot is a platform-neutral orchestrator-sync discipline reusable by ANY project whose conductor depends on a long-running autonomous test/QA framework; the consuming project supplies the concrete channel paths and payload schema per §11.4.35.

Composes with §11.4.4 (test-interrupt-on-discovery — the live stream is what lets the conductor STOP the instant a fresh defect's FAIL event lands), §11.4.5 / §11.4.69 (every verdict event cites its captured-evidence path), §11.4.27 (no-fakes — the framework exercises the real system, the channel reports real verdicts), §11.4.28 (owned-submodule decoupling — channel + payloads stay project-agnostic), §11.4.45 (status vocabulary + sink-side evidence reuse), §11.4.52 (autonomous-validation — the channel is how an autonomous session reports without a human watching the screen), §11.4.89 (background test execution — the channel is the conductor's window into a backgrounded long test), §11.4.94 / §11.4.97 (zero-idle — tailing the live channel replaces blind polling/idle waiting).

Propagation. Propagation gate `CM-COVENANT-114-116-PROPAGATION` enforces the literal anchor `11.4.116` across the consumer fleet; paired §1.1 meta-test mutation strips the literal → the gate FAILs. Recommended per-family gate `CM-AUTONOMOUS-FRAMEWORK-SYNC-CHANNEL` (the framework emits an append-only event stream + an atomically-rewritten status snapshot, and every verdict event carries an evidence path); paired §1.1 mutation strips the evidence-path field from a verdict event → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md §11.4.116 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--no-sync-channel`, `--end-of-session-report-suffices`, `--verdict-without-evidence-path`, `--torn-status-write-OK` flag exists.

§ . . – Computer-vision / OCR pixel-oracle fallback for non-introspectable UIs
mandate (. . -dev remediation,
– –)

Short tag: `cv-ocr-pixel-oracle-fallback` .

Forensic anchor (genericised, 2026-06-03). Several apps under test (TV-Compose / leanback streaming apps, canvas-rendered players, games) return a near-empty accessibility / semantic tree — `uiautomator dump` yields a hierarchy with zero clickable nodes, so a hierarchy-driven test cannot find the play tile, cannot read the on-screen caption, cannot confirm what the user sees. The §11.4.52 rule already SKIPS UI-driven when

`uiautomator dump | grep -c clickable=true == 0` . The remediation's lesson: SKIP is not the only answer — when the semantic tree is blank, the test MUST fall back to driving + asserting via the PIXELS (computer-vision template-match to locate and tap controls; ROI OCR to read on-screen text), because the pixels are what the end user actually experiences, and a hierarchy-only tool is NOT a content oracle.

The mandate. Any test that needs to drive a UI control OR assert on-screen content MUST NOT assume the accessibility / semantic / DOM hierarchy is the source of truth for what the user sees. When the hierarchy is blank, partial, or known-unreliable for the app under test (TV-Compose, leanback, canvas/`SurfaceView` /GL, games, custom-rendered UIs), the test MUST fall back to a **pixel oracle**: (1) **drive** input by computer-vision template-match (locate a control by its rendered appearance, then tap its screen coordinates) rather than by a hierarchy node id; (2) **assert** on-screen content by ROI OCR (read the actual rendered text in the region the user reads — caption strip, title, error overlay) with a per-word confidence floor + region-of-interest per §11.4.107(12), never by a hierarchy text attribute that may be absent or stale. The test tool MUST be chosen for its ability to BOTH drive input AND read pixels — a hierarchy-only tool cannot serve as a content oracle for these UIs. The pixel path is the §11.4.52 "near-empty hierarchy proves UI-driven INFEASIBLE" fallback made constructive: not SKIP, but pixel-drive.

Anti-bluff for the pixel oracle (§11.4.107). The CV/OCR analyzer is itself a potential bluff surface and MUST be self-validated per §11.4.107(10): a golden-good fixture (the control present / the expected caption rendered) MUST PASS and

a golden-bad fixture (control absent / wrong-or-empty caption) MUST FAIL, wired into the meta-test. Thresholds (match confidence, OCR confidence floor) MUST be calibrated on the project's own captured frames, not hardcoded from literature (§11.4.107(13) / §11.4.6). An OCR PASS that cannot also FAIL its golden-bad fixture is a §11.4 bluff gate.

Honest boundary (§11.4.6). When BOTH the hierarchy is blank AND the pixel oracle is infeasible (e.g. the target output is a protected/secure surface that captures black per §11.4.112, or the content is geo-unreachable), the test SKIPS-with-reason per §11.4.3 (`topology_unsupported` / `operator_attended` / `geo_restricted`), never a fake PASS, and the operator-attended path is a tracked migration item per §11.4.52, not the permanent answer.

Classification: universal (§11.4.17) — driving + asserting via pixels when the semantic tree is blank is a platform-neutral test-oracle discipline reusable by ANY project testing custom-rendered / canvas / game / Compose / leanback UIs; the consuming project supplies its concrete CV/OCR toolchain (template-match library, OCR engine), capture mechanism, and calibrated thresholds per §11.4.35.

Composes with §11.4.5 (captured-evidence quality — OCR overlay census), §11.4.6 (no-guessing — thresholds calibrated, not assumed), §11.4.48 (UI-driven — pixel-drive is the fallback when the hierarchy path is infeasible), §11.4.49 (dual-approach — pixel-drive UI variant complements the intent variant), §11.4.52 (autonomous-validation — strict refinement: near-empty hierarchy → pixel-drive, not only SKIP), §11.4.107 (AV/test-validation — the OCR/CV analyzer is self-validated golden-good/golden-bad, ROI + confidence floor; §11.4.117 is its UI-driving companion), §11.4.112 (structural-impossibility — secure-surface black-capture is the honest boundary where even pixel-drive cannot read content).

Propagation. Propagation gate `CM-COVENANT-114-117-PROPAGATION` enforces the literal anchor `11.4.117` across the consumer fleet; paired §1.1 meta-test mutation strips the literal → the gate FAILs. Recommended per-family gate `CM-CV-OCR-PIXEL-ORACLE-FALLBACK` (a test targeting a known-blank-hierarchy app drives + asserts via a self-validated CV/OCR pixel oracle, not a hierarchy-text attribute); paired §1.1 mutation makes the pixel-OCR analyzer PASS its golden-bad fixture → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md §11.4.117 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no

```
--hierarchy-is-content-oracle , --skip-pixel-fallback ,
--unvalidated-ocr-OK , --hardcoded-ocr-threshold-OK flag exists.
```

```
$ . . - Discovery-pressure to
confirm known-issue-set completeness mandate
( . . -dev remediation, - - )
```

Short tag: `discovery-pressure-completeness` .

Forensic anchor (genericised, 2026-06-03). A remediation cycle has a list of reported defects. Fixing exactly those is NECESSARY but NOT SUFFICIENT — the reported set is what the operator happened to notice, biased by what they tested ("we only see what we test"). The remediation ran parallel discovery + stress streams across ALL target devices, exercising subsystems the reported defects did NOT touch, to CONFIRM the reported set IS the complete critical set (and to surface any unreported defect before the user does). The discovery is anti-bluff only if it provides PROVABLE coverage: a list of the subsystems / journeys actually exercised, each with the outcome (no-new-issue vs new-issue-filed), so "we found nothing else" is backed by "here is everything we looked at," not by "we stopped looking."

The mandate. A remediation / release cycle **MUST NOT** treat "every reported defect is fixed" as "the build is good." After (or in parallel with) fixing the reported set, the cycle **MUST** run a discovery + stress pass across **ALL** target devices / environments that deliberately exercises subsystems, journeys, and edge cases **BEYOND** the reported defects — to confirm the reported set is the **COMPLETE** critical set and to surface unreported defects before the end user does. The discovery pass **MUST** produce **PROVABLE** coverage: an enumerated list of the subsystems / user-journeys / stress scenarios actually exercised, each with its outcome (no-new-issue, or a new tracker entry filed per §11.4.15 / §11.4.16). "We found no other issues" is a §11.4 bluff unless accompanied by "here is the

enumerated set we exercised" — absence of evidence of looking is not evidence of absence. Newly-discovered defects trigger §11.4.4 test-interrupt + the §11.4.114/ §11.4.115 isolation → RED → fix loop.

Why necessary-not-sufficient. A cycle that fixes only the reported defects ships with the operator's coverage bias baked in: the subsystems the operator did not test on the broken build are still unverified on the fixed build. Discovery pressure (everyday-user-simulation journeys + chaos/stress per §11.4.85, run autonomously per §11.4.52 across every device per §11.4.119) is how the cycle earns the claim "this is the complete critical set," converting an unbounded "are there other bugs?" into a bounded, evidenced "here is what we exercised and what we found."

Honest boundary (§11.4.6). Discovery pressure reduces the unknown-unknown surface; it does not prove zero remaining defects (no finite test set can). The claim the cycle earns is "the reported set + the enumerated discovery set are all addressed," with the discovery set's coverage explicitly bounded — never "the build is bug-free." Subsystems NOT exercised are stated as such (an honest coverage gap per §11.4.3), not silently implied clean.

Classification: universal (§11.4.17) — discovery-pressure to confirm known-issue-set completeness is a platform-neutral release-confidence discipline reusable by ANY project with a remediation/release cycle; the consuming project supplies its subsystem inventory, everyday-user-journey set, stress scenarios, and target-environment matrix per §11.4.35.

Composes with §11.4.4 (test-interrupt — a discovery finding STOPS the cycle), §11.4.5 / §11.4.69 (the discovery pass captures evidence per subsystem exercised), §11.4.25 (full-automation-coverage — the discovery set is part of the coverage ledger), §11.4.40 (full-suite retest — discovery runs alongside / inside it), §11.4.42 (iteration-discipline — discovery is the smoke-then-full progression's breadth dimension), §11.4.52 (autonomous-validation — discovery journeys run without a human), §11.4.85 (stress + chaos — discovery pressure IS the stress dimension applied for completeness), §11.4.114 / §11.4.115 (a discovered defect enters the isolation → RED → fix loop), §11.4.119 (single-resource-owner — parallel per-device discovery needs the ownership partition).

Propagation. Propagation gate `CM-COVENANT-114-118-PROPAGATION` enforces the literal anchor `11.4.118` across the consumer fleet; paired §1.1 meta-test mutation strips the literal → the gate FAILs. Recommended per-family gate `CM-`

DISCOVERY - COVERAGE - ENUMERATED (a remediation/release cycle records an enumerated subsystem/journey discovery-coverage list with per-item outcome, not a bare "no other issues found"); paired §1.1 mutation replaces the enumerated list with a bare no-issues claim → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md §11.4.118 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no

```
--reported-set-suffices , --skip-discovery-pass ,  
-i no-issues-without-coverage-list , --assume-complete flag exists.
```

1 1 8
2026 06 03

§ . . — Single-resource-owner
partitioning for parallel hardware testing
mandate (. . -dev remediation,
- -)

Short tag: single-resource-owner-partitioning .

Forensic anchor (genericised, 2026-06-03). Multiple parallel test/discovery streams ran against shared physical hardware (devices D3 + D4, an HDMI audio sink). Two streams driving media playback on the SAME device at once corrupt each other's evidence — overlapping audio routes the sink reports the wrong codec, concurrent `am start` fights over the foreground, a second stream's input events land in the first stream's app. The remediation enforced: exactly ONE stream OWNS each device's exclusive/media resource at a time; other streams targeting that device are read-only (probes / dumphsys / log capture). Parallelism is partitioned BY DEVICE — stream A owns D3's media, stream B owns D4's media — so the parallel-development pipeline (§11.4.58 / §11.4.103) runs without the streams clobbering each other's captured evidence.

The mandate. When multiple parallel work/test/discovery streams exercise SHARED hardware or any exclusive-access resource (a media-playback path, a single HDMI/audio sink, an exclusive device handle, a single serial/JTAG line, a GPU under exclusive capture), exactly ONE stream MUST own each such resource at a time. The exclusive owner drives the resource (playback, input injection, capture); every other concurrent stream targeting the same resource MUST be

READ-ONLY (passive probes — `dumpsys` / `/proc` / `/sys` reads / sink-side network probes / log tails) for the duration. Parallelism MUST be partitioned by resource: distinct devices / sinks / handles run fully concurrent (stream-per-device), but the same device's exclusive resource is single-owner. The ownership MUST be enforced by an advisory lock / token (per §11.4.58 L3 contention-path locking + §11.4.84 working-tree quiescence's hardware analogue), not by convention — a second stream MUST be unable to silently start media on a device another stream owns. Ownership is event-driven: a stream claims a device when its exclusive resource frees and releases it the moment its work completes, so the next queued stream can claim it (per §11.4.103 auto-backfill).

Why single-owner. Concurrent drivers of one exclusive resource produce CROSS-CONTAMINATED evidence: the sink reports whichever stream's audio won the race, the foreground belongs to whichever `am start` landed last, input events interleave. A PASS captured under contention is a §11.4 evidence-integrity bluff — the captured artefact does not reflect what the test under inspection actually did. Single-owner partitioning is the precondition that makes per-stream captured evidence trustworthy under parallelism.

Honest boundary (§11.4.6). A read-only probe stream is genuinely read-only — it MUST NOT issue any state-changing command (no `am start`, no input injection, no `setprop`, no playback) against a device it does not own; a "read-only" stream that mutates state is the contention the partition forbids. When the resource genuinely cannot be partitioned (a single device the whole batch needs to drive), the streams MUST serialize on it (single-owner over time), not run concurrently and hope.

Classification: universal (§11.4.17) — single-owner-per-exclusive-resource partitioning is a platform-neutral parallel-testing discipline reusable by ANY project with parallel streams contending over shared hardware / exclusive handles; the consuming project supplies its resource inventory (which devices / sinks / handles are exclusive) and lock mechanism per §11.4.35. Strict refinement of §11.4.58 / §11.4.103 for the hardware-contention case.

Composes with §11.4.5 / §11.4.69 (captured-evidence integrity — single-owner is the precondition for trustworthy per-stream evidence), §11.4.13 (sink-side evidence — a shared sink is the canonical single-owner resource), §11.4.50 (deterministic consistency — contended evidence is non-deterministic, defeating N-iteration stability), §11.4.58 (parallel-development PWU — §11.4.119 is the hardware-resource refinement of its L3 contention-path locking), §11.4.82(F) (parallel multi-

device testing — per-device streams ARE the partition), §11.4.84 (working-tree quiescence — the hardware analogue: a stream owns the device's exclusive state cleanly before it drives it), §11.4.103 (continuous parallel streams — single-owner + event-driven release enables the auto-backfill), §11.4.118 (discovery pressure — parallel per-device discovery needs this partition).

Propagation. Propagation gate `CM-COVENANT-114-119-PROPAGATION` enforces the literal anchor `11.4.119` across the consumer fleet; paired §1.1 meta-test mutation strips the literal → the gate FAILs. Recommended per-family gate `CM-SINGLE-RESOURCE-OWNER-PARTITION` (parallel hardware-test streams acquire a per-resource ownership lock before any state-changing command; concurrent same-resource streams are read-only); paired §1.1 mutation removes the ownership-lock acquisition (letting two streams drive one device) → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md §11.4.119 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--allow-concurrent-resource-drivers`, `--skip-ownership-lock`, `--read-only+may-mutate`, `--contended-evidence-OK` flag exists.

1 1 8

2026-06-03

§ . . — Fix-breaks-its-own-gate
reconciliation mandate (. . -dev
remediation, - -)

Short tag: `gate-reconciliation-not-fake-pass` .

Forensic anchor (genericised, 2026-06-03). A correct fix removed the regression-causing behaviour (the audio fix removed the `handledByDualHdmi` short-circuit that diverted 5.1 off the multichannel thread). But an EARLIER gate had been authored to assert the OLD — now-known-broken — behaviour's presence. After the fix, that gate FAILED. The FAIL was the CORRECT signal that the fix worked: the gate was asserting a behaviour the fix deliberately removed. Two wrong responses are forbidden — fake-passing the gate (editing it to always pass), or

reverting the correct fix to satisfy the stale gate. The right response is to RECONCILE the gate: rewrite it to assert the NEW mechanism (evidence-backed), so it once again guards the now-correct behaviour.

The mandate. When a correct fix causes a pre-existing gate / test to FAIL because that gate asserted the OLD (now-removed-or-changed) behaviour, the gate FAIL is the CORRECT signal that the fix landed — it MUST NOT be suppressed by either of the two forbidden responses: (1) **fake-passing** the gate — editing it to `always pass`, weakening its assertion to a tautology, or deleting it to make the FAIL disappear (a §11.4 bluff at the gate layer, and a §11.4.84 mutation-residue risk); (2) **reverting the correct fix** to satisfy the stale gate (re-introducing the defect to keep a now-wrong gate green). The required response is **RECONCILIATION**: rewrite the gate so it asserts the NEW mechanism the fix introduced, backed by captured evidence of the new correct behaviour, so the gate once again guards the right invariant. After reconciliation, the gate's paired §1.1 mutation MUST be updated too (the mutation must break the NEW invariant). The reconciliation MUST be a visible, evidence-cited change — "rewrote gate X to assert because fix Y removed ; new captured evidence: " — never a silent assertion-weakening.

Distinguish reconciliation from bluffing. Reconciliation rewrites the gate to assert a REAL new invariant (the fix's intended mechanism), proven by captured evidence, with the paired mutation updated to break it. Bluffing weakens the gate so it can no longer FAIL on the defect class it was meant to catch. The discriminator: after reconciliation the gate + its mutation still form a valid §1.1 pair (mutate the new invariant → gate FAILs); a bluffed gate's mutation no longer makes it FAIL (the assertion became a tautology). A reconciled gate is provably still a guard; a bluffed gate is provably blind.

Honest boundary (§11.4.6). A gate FAIL after a fix is NOT automatically "the gate is stale, reconcile it." It MUST first be investigated per §11.4.102 — the FAIL may be the gate correctly catching a REGRESSION the fix introduced (the fix broke something else), in which case the FIX is wrong, not the gate. Reconcile ONLY when investigation PROVES the gate asserted old-correct-now-removed behaviour AND the new behaviour is the intended, evidence-confirmed mechanism. "The gate is just stale" without that investigation is a §11.4.6 guess that can mask a real regression.

Classification: universal (§11.4.17) — reconcile-the-gate-don't-fake-pass-and-don't-revert is a platform-neutral fix-vs-gate-conflict discipline reusable by ANY project with regression gates; the consuming project supplies its gate framework and paired-mutation mechanism per §11.4.35.

Composes with §11.4.1 (FAIL-bluffs — fake-passing the gate is the inverse bluff), §11.4.4 (test-interrupt — a post-fix gate FAIL triggers investigation before reconcile), §11.4.6 (no-guessing — "stale gate" is a hypothesis until investigation proves it), §11.4.84 (working-tree quiescence — a gate edited to `always pass` is mutation residue that must never ship), §11.4.102 (systematic-debugging — the mandatory investigation that distinguishes stale-gate from new-regression), §11.4.108 (four-layer fix-verification — the reconciled gate asserts the fix's runtime signature / new mechanism), §1.1 (paired mutation — the reconciled gate's mutation must break the NEW invariant, re-proving it is a guard not a tautology).

Propagation. Propagation gate `CM-COVENANT-114-120-PROPAGATION` enforces the literal anchor `11.4.120` across the consumer fleet; paired §1.1 meta-test mutation strips the literal → the gate FAILs. Recommended per-family gate `CM-GATE-RECONCILED-NOT-FAKE-PASSED` (when a fix changes a behaviour a gate asserted, the gate's reconciliation is an evidence-cited rewrite whose paired mutation still makes it FAIL — never an assertion weakened to a tautology); paired §1.1 mutation weakens a reconciled gate to `always pass` and asserts its mutation no longer makes it FAIL → the meta-gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

GATE SEAM-PLACEMENT — a gate whose precondition the gated work itself produces is on the WRONG SEAM (extension, research-derived, 2026-07-17). §11.4.120's trigger is a gate that FAILs — i.e. a gate that RAN. Its sibling case is a gate that CANNOT RUN in the ordering imposed on it, and the required response is different: not reconciliation, but re-seating the gate at the seam where it is satisfiable. Forensic FACT (genericised): work was ordered "run checks E1-E4, then implement P1 **only if** E1-E4 pass" — but the gate's own text defined E2 as running **on a P1-gated build**, so E2 could not exist before P1 existed; the order was unexecutable by construction. The source of truth, read exactly, gated something else entirely: building behind a default-off flag was permitted, and it was **ENABLING** the flag that required E1-E4 — the gate was on ACTIVATION, and had been imposed on AUTHORSHIP. Therefore: **(a) EXECUTABILITY IS CHECKED BEFORE A GATE IS IMPOSED** — if gate G requires artifact A, and A is produced

by the very work G gates, then G is unsatisfiable at that seam and belongs at a LATER one (activation / enablement / release), NEVER at authorship; imposing it at authorship makes the work impossible to start and the gate impossible to pass, which is not rigor, it is a deadlock. **(b) RE-READ THE SOURCE OF TRUTH**

BEFORE RE-SEATING — the circularity is usually an ARTIFACT OF THE ORDER, not of the gate: the gate's own text names its real seam, and that text governs over any restatement of it (a restatement that moved the gate from "enabling" to "writing" is a §11.4.6 misstatement of the requirement, and the fix is to obey the source, not to weaken the gate). **(c) THE FORBIDDEN RESPONSES ARE §11.4.120's OWN**

— an unsatisfiable gate MUST NOT be fake-passed, deleted, weakened to a tautology, nor improvised past by building the artifact and retro-declaring the gate met; and the gated work MUST NOT be abandoned as "blocked" when the gate simply sits on the wrong seam. **Honest boundary (§11.4.6):** an unrunnable gate is NOT automatically mis-seated — investigate per §11.4.102 first; it may be correctly seated and its precondition genuinely missing, which is the §11.4.69

`artifact_not_yet_built` NOT-YET-RUNNABLE state (§11.4.135 seam-dependent), not a seam defect. Reporting the circularity is mandatory and already required by §11.4.6 / §11.4.66 / §11.4.101 / §11.4.201(4) — this clause adds the seam-placement rule those anchors do not carry. §11.4.110's clash detector cannot catch this class: its input is a change DIFF, and an ordering instruction is not a diff.

Gate contract for the seam-placement clause (extension, 2026-07-17). `CM-GATE-RECONCILED-NOT-FAKE-PASSED` is EXTENDED: an imposed gate whose required precondition artifact is produced by the very work it gates → FAIL (the gate is mis-seated and MUST be re-seated at activation/enablement/release), and the FAIL names the circularity rather than the work. Paired §1.1 mutation: impose an authorship-seam gate whose precondition only the gated work can produce → the gate FAILs; re-seat it at activation → PASSes. Its golden-FALSE fixture MUST include a correctly-seated gate whose precondition is merely not-yet-built (the §11.4.69 `artifact_not_yet_built` state) — the gate MUST NOT fire on that, or it becomes the §11.4.201(1) false positive that condemns a healthy gate. Gate-code = a separate work item; this is the contract, not a claim the code has shipped (§11.4.6).

Canonical authority: this Constitution.md §11.4.120 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--fake-pass-stale-gate`, `--revert-fix-for-gate`, `--weaken-assertion-to-pass`, `--delete-failing-gate`, `--impose-unsatisfiable-gate`, `--improvise-past-circular-order`, `--gate-authorship-on-its-own-output` flag exists.

1 1 8

2026 06 03

§ . . — No-commit-while-build-writes-tracked-artifacts mandate (. . -dev remediation, - -)

Short tag: `no-commit-during-artifact-write` .

Forensic anchor (genericised, 2026-06-03). The build pipeline's per-component build steps write produced artifacts (forked-app APKs) INTO tracked directories that are under version control (`prebuilt_apps/`). Running `git add -A /` a commit while those steps are mid-write stages a PARTIAL or stale artifact — a race between the build writing the file and git reading it — so the commit captures a half-written or pre-rebuild artifact. The remediation deferred the commit to BUILD-COMPLETION: the tree is quiescent (no build step writing tracked files) AND the freshly-built artifacts are captured whole (no stale prebuilt in git). This is the artifact-write analogue of §11.4.84 (no commit while a mutation gate is in flight) at the build-output layer.

The mandate. A commit (especially `git add -A /` any broad stage) MUST NOT run while a build / packaging / generation step is actively writing artifacts INTO tracked directories. Doing so races the writer and stages a partial or stale artifact — the commit captures bytes that are mid-write or pre-rebuild, a §11.4.108 SOURCE → ARTIFACT integrity failure landed in version control. The commit MUST be deferred until the build step that writes tracked artifacts has COMPLETED, so that: (1) the working tree is quiescent at the artifact layer (no step writing a tracked file); (2) the committed artifacts are the FRESH, whole outputs of the completed build (no stale pre-rebuild artifact, no half-written file). Before committing tracked build outputs, the committer MUST verify the writing step finished (process exit, completion marker, or per-artifact mtime ≥ build-start) — a build still in flight that writes tracked dirs is a HOLD on the commit, not a race to win.

Composition with §11.4.84. §11.4.84 forbids committing while a mutation gate is in flight in the checkout (mutation residue). §11.4.121 is the BUILD-OUTPUT analogue: forbid committing while a build is in flight WRITING TRACKED ARTIFACTS (partial/stale artifact residue). Both close the same class of defect — a commit captures transient, non-final tree state — at two different write-sources (mutation experiments vs build outputs). Where build outputs land OUTSIDE version control (gitignored `out/` / `dist/`), the race does not apply; the mandate binds specifically to build steps writing into TRACKED directories.

Honest boundary (§11.4.6). "The build probably finished" is not "the build finished" — verify with a completion signal (exit code / marker / mtime), not an assumption, before committing tracked artifacts. Committing source changes that DON'T touch the build's tracked-artifact directories is fine mid-build (those files are not being written by the build); the hold applies only to the tracked dirs the in-flight build writes.

Classification: universal (§11.4.17) — deferring commit until a tracked-artifact-writing build completes is a platform-neutral commit-integrity discipline reusable by ANY project whose build writes produced artifacts into version-controlled directories; the consuming project supplies which build steps write which tracked directories and its completion-signal mechanism per §11.4.35. The PROJECT-SPECIFIC instance (which tracked directory — e.g. a prebuilt-apps directory — and which build steps write it) is recorded in the consuming project's own governance per §11.4.35.

Composes with §11.4.6 (no-guessing — verify build completion, never assume), §11.4.30 (no-versioned-build-artifacts — artifacts that ARE legitimately tracked, like signed prebuilts, are the exact case this mandate protects from partial commit; gitignored build outputs are exempt), §11.4.58 / §11.4.103 (parallel-development — the commit/merge stage is serialized AFTER the build stage for this reason), §11.4.84 (working-tree quiescence — §11.4.121 is its build-output analogue), §11.4.88 (background-push — commit timing is decoupled from push, but the commit itself still waits for artifact-write completion), §11.4.108 (four-layer fix-verification — a partial/stale committed artifact is the ARTIFACT-layer integrity failure this prevents from entering version control), §11.4.96 (safe-parallel-work-with-long-build — committing tracked build outputs mid-build is an UNSAFE operation in its catalogue).

Propagation. Propagation gate `CM-COVENANT-114-121-PROPAGATION` enforces the literal anchor `11.4.121` across the consumer fleet; paired §1.1 meta-test mutation strips the literal → the gate FAILs. Recommended per-family gate `CM-NO-COMMIT-DURING-ARTIFACT-WRITE` (the commit wrapper refuses to stage tracked build-artifact directories while a build step writing them is in flight — checks a build-in-progress lock / completion marker); paired §1.1 mutation removes the build-in-flight check from the commit wrapper → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md §11.4.121 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--commit-during-build`, `--stage-partial-artifact-OK`, `--assume-build-finished`, `--skip-build-completion-check` flag exists.

2026 06 03

§ . . — No-silent-removal-of-existing-components-without-operator-confirmation mandate (User mandate, — —)

Forensic anchor — verbatim user mandate (2026-06-03):

"Never ever remove any application, system component or service from already existing codebase / System without interactively asked question to us! THIS IS MANDATORY RULE / CONSTRAINT!"

Forensic case study (this project, FACT — captured). During the 1.1.8-dev burn-down, two shipped capabilities — F2 (an Apple-TV-class application) and F4 (a Huawei HMS / Mobile-Services component) — were removed from the existing System WITHOUT first asking the operator. The operator reversed both removals. A removal that the operator has to discover and reverse after the fact is a defect of the same severity class as a §11.4 PASS-bluff: the System silently lost a user-facing capability the operator never agreed to drop.

The mandate. No application, system component, service, package, feature, driver, module, library, prebuilt asset, or any other already-existing capability of the existing codebase / shipped System may be removed (deleted, dropped from the package set, disabled-into-non-shipping, un-bundled, de-listed, or otherwise made unavailable to the end user) WITHOUT FIRST interactively asking the operator and receiving an EXPLICIT keep-or-remove decision. The question MUST be posed through the platform's interactive clarification mechanism per §11.4.66 (`AskUserQuestion` on Claude Code; the documented interactive prompt on other runtimes) — NEVER a free-text "should I remove X?" buried in narrative, NEVER a silent removal justified post-hoc, NEVER an autonomous removal decision. A silent removal is a **release blocker** regardless of how well-intentioned the rationale (deduplication, "it was broken anyway", "geo-restricted", "incompatible", "superseded") — the operator decides, the agent asks.

What counts as a removal (non-exhaustive). Deleting an app/APK/binary from the build's package set (e.g. a project's `PRODUCT_PACKAGES` / `device.mk` / equivalent manifest), removing a service from the init/boot/service-registry set, dropping a kernel module / driver / config from the shipping configuration, un-bundling a prebuilt asset, deleting a submodule or a submodule's shipped output, removing a feature flag that gated a live capability, or any edit whose NET EFFECT is "an end-user capability that shipped before no longer ships." Adding, replacing-with-equivalent-or-better (where the operator has approved the replacement), or fixing a capability is NOT a removal. When uncertain whether an edit constitutes a removal, treat it AS a removal and ask (per §11.4.6 no-guessing + §11.4.101 block-only-when-truly-needed — removal of an existing user-facing capability is high-blast-radius and MUST be operator-confirmed, never autonomously decided).

Composition. §11.4.122 composes with §11.4.66 (interactive-clarification — the removal question is asked via the platform interactive mechanism, not free-text), §11.4.101 (autonomous-decision-over-blocking — a removal of an existing user-facing capability is exactly the irreversible-ish, high-blast-radius class §11.4.101 RESERVES for an operator question rather than an autonomous decision; §11.4.122 makes the "ask the operator" outcome mandatory for this class, NOT optional), §11.4.90 (Obsolete classification — the tracked path to actually DROP a capability is: ask the operator → operator approves → mark the item `obsolete` (→ `Fixed.md`) with `Obsolete-Details` reason `feature-removed` AND an operator-approval citation → then remove; the removal never precedes the operator's yes), §11.4.112 (structural-impossibility won't-fix — even a capability

proven structurally impossible to keep is closed via the §11.4.90 path with operator awareness, not silently deleted), §11.4.6 (no-guessing — "the operator probably wants this gone" is the exact guess forbidden), §11.4.40 / §11.4.42 (full-suite-retest / iteration-discipline — a removal landed without the operator's yes fails the release gate).

Enforcement. Propagation gate `CM-COVENANT-114-122-PROPAGATION` enforces the literal anchor `11.4.122` across the consumer fleet (every `CLAUDE.md` / `AGENTS.md` / `QWEN.md`). Recommended gate `CM-NO-SILENT-COMPONENT-REMOVAL` (a diff-driven detector that flags a removal of a shipped capability — a deletion from the project's package/service/module manifest set — without a matching operator-approval citation in the commit / tracker entry, per the §11.4.110 change-impact-clash pattern; a removal whose tracker item is `Obsolete` (→ `Fixed.md`) with reason `feature-removed` + operator-approval citation passes). Paired §1.1 meta-test mutations: strip the `11.4.122` literal → propagation gate FAILs; introduce a manifest removal with no operator-approval citation → `CM-NO-SILENT-COMPONENT-REMOVAL` FAILs. (Gate-code implementation lands as a separate follow-up work item; this anchor defines the contract.)

Classification: universal (§11.4.17) — "never silently remove an existing end-user capability without operator confirmation" is a platform-neutral discipline reusable by ANY project that ships a set of user-facing capabilities; the consuming project supplies its concrete capability-manifest paths (package set, service registry, module config) per §11.4.35.

Canonical authority: this Constitution.md §11.4.122 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no

```
11 4 123
--remove-without-asking , --silent-removal , --autonomous-removal-OK ,
--dedup-removal-exempt , --it-was-broken-anyway flag exists.
2026 06 03
```

§ . . — Rock-solid-proof-or-deep-research mandate (User mandate,
- -)

Forensic anchor — verbatim user mandate (2026-06-03):

"Every single reported issue MUST BE fully and 100% validated with rock solid proofs! Nothing can be considered fixed or completed without hard evidence! No false results or bluff(s) of any kind is allowed! If we are not sure on how to achieve full testing, validation and verification of something we MUST ALWAYS perform deep web research for all possible data (articles, documentation, guides, and other resources) and opensourced codebases which we can use to solve our problems and perform testing with validation and verification which produces rock-solid evidence(s) and leaves no space for false results or any kind of bluff!"

Forensic case study (FACT — captured). In the 1.1.8-dev remediation cycle the validation method for two feature classes was, at first, genuinely unclear: relocating a `FLAG_SECURE` secure surface to a secondary display (whose pixel capture returns black) and asserting on-screen content in non-introspectable streaming-app UIs (whose accessibility hierarchy is blank). Rather than declaring those classes "untestable" or accepting a metadata-only PASS, the cycle performed deep web research (`docs/research/testing_frameworks_20260603/`) into platform/HDCP behaviour, computer-vision template-matching, ROI OCR with confidence floors, freeze-detection / frame-advance liveness oracles, and network-side sink probes. That research yielded the CV/OCR/liveness/sink-probe oracle stack (now §11.4.107 + §11.4.112 + §11.4.117) that made rock-solid, evidence-producing validation possible where it had appeared impossible. The lesson — "unclear how to validate" is a research trigger, NEVER a bluff licence — is now constitutional.

The mandate. Every single reported issue, every fix, and every claimed completion MUST be fully and 100% validated with rock-solid CAPTURED proof per §11.4.5 / §11.4.69 / §11.4.107 before it may be marked fixed / implemented / completed (the §11.4.33 terminal closure vocabulary). Hard, captured evidence is the bar — nothing may be considered fixed or complete without it. Metadata-only PASS, configuration-only PASS, absence-of-error PASS, and grep-without-runtime PASS are all forbidden (§11.4 / §11.4.1) — no false results and no bluff of any kind, at any layer, are allowed.

The research-or-don't-bluff rule (the operative addition). When the agent is UNSURE how to fully test, validate, or verify something — when no obvious evidence-producing method exists, OR the candidate method would yield only

metadata/config/absence-of-error evidence (and is therefore a bluff per §11.4.1) — the agent MUST ALWAYS first perform deep web research per §11.4.8 + §11.4.99 (official documentation, technical articles, guides, vendor references, standards, issue trackers, and open-source codebases the project can reuse or adapt) to DISCOVER or BUILD a validation method that produces rock-solid evidence and leaves no space for a false result. Declaring something "untestable", "not automatable", or accepting a metadata-only PASS WITHOUT having first exhausted this deep-research path is itself a §11.4.123 violation — the same severity class as a PASS-bluff. The research output (cited source URLs + the evidence-producing method it yielded, OR the literal "NO external solution found — original work" per §11.4.8) is the captured proof that the path was exhausted. Only after that research genuinely fails to yield any evidence-producing method may the item be classified `PENDING_FORENSICS: / Operator-blocked (§11.4.21) / structurally-impossible won't-fix (§11.4.112)` — with the cited research as the evidence that the classification is earned, never a convenience.

Composition. §11.4.123 composes with §11.4.5 (captured-evidence quality is the proof shape), §11.4.6 (no-guessing — "probably fixed" without captured evidence is forbidden; the closure must be FACT), §11.4.8 (deep-web-research-before-implementation — §11.4.123 extends it from "before designing a fix" to "whenever the VALIDATION method is unclear"), §11.4.52 (autonomous-validation — the research must yield an autonomous, evidence-producing path wherever one is discoverable), §11.4.69 (universal sink-side positive-evidence taxonomy — every PASS cites a taxonomy-class artefact), §11.4.99 (latest-source verification of the discovered method), §11.4.107 (the AV/test-validation oracle techniques the 1.1.8-dev research yielded), §11.4.118 (discovery-pressure — the same "absence of evidence of looking is not evidence of absence" principle applied to per-issue validation). It also composes with §11.4.21 / §11.4.112 (the earned escape valves — reachable ONLY after the deep-research path is exhausted with cited evidence).

Enforcement. Propagation gate `CM-COVENANT-114-123-PROPAGATION` enforces the literal anchor `11.4.123` across the consumer fleet (every CLAUDE.md / AGENTS.md / QWEN.md). Recommended gate `CM-ROCK-SOLID-PROOF-OR-RESEARCH` (a closure-audit detector that flags any item closed `Fixed / Implemented / Completed` whose tracker entry lacks a captured-evidence artefact path per §11.4.69, AND any `PENDING_FORENSICS: / Operator-blocked / structurally-impossible` classification whose entry lacks a cited deep-research trail per §11.4.8/§11.4.99). Paired §1.1 meta-test mutations: strip the

11.4.123 literal → propagation gate FAILs; close an item with a metadata-only PASS and no captured-evidence path → CM-ROCK-SOLID-PROOF-OR-RESEARCH FAILs; classify an item `Operator-blocked` with no cited research trail → the same gate FAILs. (Gate-code implementation lands as a separate follow-up work item; this anchor defines the contract.)

Classification: universal (§11.4.17) — "every closure carries rock-solid captured proof, and unclear-how-to-validate triggers deep research rather than a bluff" is a platform-neutral discipline reusable by ANY project; the consuming project supplies its concrete capture mechanisms + research corpora per §11.4.35.

Canonical authority: this Constitution.md §11.4.123 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no

11 4 124
`--metadata-pass-suffices` , `--skip-proof` , `--untestable-without-research` , `--config-only-closure-OK` , `--bluff-when-unsure` flag exists.
2026 06 04

§ . . — **Dead/unwired-code**
investigate-before-remove mandate (User mandate,
— —)

Forensic anchor — verbatim user mandate (2026-06-04):

"Before removing any seemingly-dead (zero-importer / unwired) codebase, we MUST investigate via git history where/how it was originally used and how it became dead. Removal is permitted ONLY when we have captured PROOF it is genuinely no longer needed — and that removal MUST be its own separate commit with a proper descriptive message. If there is no such proof, the code MUST be investigated for where/how it should be wired in properly, and any missing or unwired tests MUST be added. We MUST ALWAYS be extra careful with any codebase removal."

The defect class. "This has zero importers / is never called / is unwired, so it is dead — delete it" is a guess (§11.4.6), not a finding. A `git grep` / static-analysis "no references" result proves only *current* non-reference; it does NOT prove the code is genuinely unneeded. Code becomes apparently-dead for several distinct

reasons — its call-site was deleted by mistake (a regression: the code SHOULD still be wired in), its wiring was never finished (it is intended but unwired — the fix is to WIRE it, plus its tests), it is reflectively / dynamically / build-tag / plugin / DI-container referenced (the static tool simply cannot see the reference), or it is genuinely obsolete and safe to remove. Deleting it without first distinguishing which case applies risks silently removing a needed-but-temporarily-unwired capability — the same severity class as the §11.4.122 silent-removal defect and a §11.4 PASS-bluff at the maintenance layer.

The mandate. Before removing ANY seemingly-dead codebase element (a zero-importer / never-called / unwired / unreferenced function, method, type, file, module, package, asset, config, or build target), the agent MUST FIRST investigate via git history (`git log --follow` , `git log -S<symbol> / -G<regex>` `pickaxe` across all history, `git log --diff-filter` , `blame` on the deletion of the former call-site) to establish, as captured FACT (§11.4.6): (1) WHERE and HOW it was originally used / wired in; (2) WHEN and HOW it became dead — its call-site removed deliberately, removed by mistake (regression), never completed, or made unreachable by a refactor; (3) whether the "no references" result is real OR an artefact of a reference the static tool cannot see (reflection / dynamic dispatch / build-tags / codegen / DI / plugin registry / cross-language FFI / config-driven wiring). The investigation output (the cited commits + the determination) is the captured evidence.

Removal is conditional — proof-of-genuinely-unneeded REQUIRED. Removal is permitted ONLY when the investigation produces captured PROOF the element is genuinely no longer needed (its purpose is fully superseded, its former caller is intentionally and permanently gone, no hidden reference exists). When that proof exists, the removal MUST be its OWN SEPARATE COMMIT (not bundled with unrelated work — so it is independently reviewable and revertible, composing with §11.4.84 working-tree quiescence + §11.4.92 multi-pass) carrying a proper descriptive message that cites the git-history investigation evidence (the commits that wired it in, the commit that orphaned it, the determination that it is unneeded) and — where the element is an end-user-facing capability — the §11.4.122 operator-confirmation. The §11.4.90 tracked path applies: classify the tracker item `Obsolete (→ Fixed.md)` with the appropriate reason.

No proof ⇒ wire it in + add the missing tests (do NOT delete). If the investigation does NOT yield proof the element is unneeded, the element MUST NOT be removed. Instead it MUST be investigated for WHERE and HOW it should be wired in properly (restore the deleted call-site if its removal was a regression per §11.4.114; complete the intended wiring if it was never finished), AND any missing or unwired tests for it MUST be added (the absence of a test is itself part of why it drifted into apparent-deadness — composes §11.4.27 100%-test-type-coverage + §11.4.43 TDD-fix + §11.4.115 RED-baseline). Apparent-deadness with no removal-proof is a wiring/coverage GAP to close, never a delete-on-sight.

Extra-caution default. The agent MUST ALWAYS be extra careful with any codebase removal: when uncertain whether removal-proof is sufficient, default to NOT removing (investigate + wire + test), per §11.4.6 no-guessing + §11.4.101 (an irreversible-ish, possibly-needed removal whose safety cannot be determined from evidence is exactly the class that blocks to the operator rather than being autonomously deleted) + §11.4.122 (removal of a live end-user capability requires operator confirmation). "Probably dead" is never sufficient; the bar is captured proof.

Composition. §11.4.124 composes with §11.4.6 (no-guessing — "zero importers therefore dead" is the exact guess forbidden; the determination must be FACT from git history), §11.4.8 (deep-research — git history IS the primary research corpus for how the code was used + how it died; external research applies when the wiring mechanism is unclear), §11.4.84 / §11.4.92 (the separate-removal-commit discipline — quiescent tree, multi-pass evaluation, independently revertible), §11.4.90 (Obsolete classification is the tracked path for a proven removal), §11.4.101 (autonomous-decision-over-blocking — a removal whose safety cannot be proven blocks rather than proceeds), §11.4.114 (last-known-good-tag isolation — a call-site deleted by mistake is a regression to repair, the same diagnostic the "became dead by mistake" case demands), §11.4.122 (no-silent-removal — §11.4.124 is its complement: §11.4.122 governs removal of KNOWN-shipped capabilities by operator confirmation, §11.4.124 governs removal of SEEMINGLY-dead code by git-history proof; an end-user-facing seemingly-dead element triggers BOTH), §11.4.27 / §11.4.43 / §11.4.115 (the missing-tests obligation when the element is wired in rather than removed).

Enforcement. Propagation gate `CM-COVENANT-114-124-PROPAGATION` enforces the literal anchor `11.4.124` across the consumer fleet (every `CLAUDE.md` / `AGENTS.md` / `QWEN.md`). Recommended gate `CM-DEAD-CODE-INVESTIGATE-BEFORE-REMOVE` (a commit-audit detector: any commit whose diff is a `NET DELETION` of a code element — `function/type/file/module/package/asset/build-target` — **MUST** be a removal-only commit whose message cites the git-history investigation evidence, **OR** the deletion is part of a tracked `obsolete` (`→ Fixed.md`) item per §11.4.90; a net-deletion commit bundled with unrelated changes **OR** lacking the investigation citation **FAILS**). Paired §1.1 meta-test mutations: strip the `11.4.124` literal `→` propagation gate **FAILS**; introduce a net-deletion commit with no investigation citation and not part of an `Obsolete` item `→` `CM-DEAD-CODE-INVESTIGATE-BEFORE-REMOVE` **FAILS**. (Gate-code implementation lands as a separate follow-up work item; this anchor defines the contract.)

Classification: universal (§11.4.17) — "investigate git history + capture proof-of-unneeded before removing seemingly-dead code, otherwise wire it in and add its tests; removal is its own descriptive commit" is a platform-neutral maintenance discipline reusable by ANY codebase; the consuming project supplies its concrete static-analysis / importer-graph tooling and its hidden-reference mechanisms (reflection, DI, build-tags, codegen) per §11.4.35.

Canonical authority: this `Constitution.md` §11.4.124 in the `HelixConstitution` submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--zero-importers-means-dead` , `--delete-unwired-on-sight` , `--skip-git-history-investigation` , `--remove-without-proof` , `--bundle-removal-with-other-work` flag exists.

2026 06 04

§ . . — **Code-review-agent gate**
before pre-build + main build (mandatory multi-layer review) (User mandate,
— —)

Forensic anchor — verbatim user mandate (2026-06-04):

"After all fixes/changes/implementations are done, BEFORE running pre-build tests and the main build, dispatch code-review agent(s) that analyze all work done + all existing data/facts + the existing codebase + current git history to determine quality, safety, and whether the fixes/changes will REALLY work; they MUST validate and verify that every test covering the fixes/changes genuinely validates the work with NO chance of false results or bluff of any kind. Any finding MUST be fixed, polished, improved, and covered with additional tests before the build proceeds. Multiple strong layers of checks."

The defect class. Running the pre-build gate sweep and the main (artifact) build directly after a batch of fixes/changes — without an independent review pass first — lets a defective fix, an error-prone change, a bluff test, or a test that does not genuinely validate the work-under-test consume an expensive build cycle and (worse) ship green. Pre-build gates and the build prove the codebase is internally consistent + buildable; they do NOT prove the work is correct, safe, will-really-work for the end user, or that the tests covering it actually catch its negation. That assurance is an independent-review responsibility that MUST run BEFORE the build, not after.

The mandate. After all fixes / changes / implementations in a batch are done, and BEFORE running the pre-build test sweep and the main build (for ANY project), the agent MUST dispatch one or more dedicated code-review agent(s) (subagent-driven by default per §11.4.70 / §11.4.20) that perform a multi-layer review. The review MUST: (1) **analyze ALL work done in the batch** — every fix / change / implementation, its source diff, and its stated intent; (2) **analyze ALL existing data + facts** — captured evidence per §11.4.5 / §11.4.69 / §11.4.107, tracker entries, prior findings, the §11.4.108 runtime-signature registry; (3) **analyze the existing codebase** — the blast radius of each change (§11.4.92 multi-pass), cross-feature interaction, contract integrity of every dependency; (4) **analyze current git history** — what each change touched, how it composes with concurrent / recent work, whether it reproduces a known-broken pattern (§11.4.114 / §11.4.124 git-history discipline); (5) **determine quality, safety, and will-it-REALLY-work** — assess whether each change is robust + not error-prone (will not solve A while creating B), safe (no host-safety / data-safety / security regression), and genuinely delivers the end-user-visible behaviour it claims (§11.4 / §107); (6) **validate + verify the tests covering the work** — confirm every test that covers a fix/change genuinely

exercises the work-under-test and catches its negation, with ZERO chance of a false result or bluff of any kind (a test that PASSES on broken-for-the-user work, a metadata-only / config-only / absence-of-error / grep-without-runtime assertion, or a gate whose paired §1.1 mutation does not make it FAIL is a finding the review MUST surface per §11.4.1 / §11.4.6 / §11.4.107 / §1.1).

Findings block the build. Any finding the review surfaces — a defect, an error-prone change, a safety risk, a will-not-really-work conclusion, a bluff/false-result-capable test, a missing-coverage gap — MUST be fixed, polished, improved, and covered with additional tests (four-layer per §11.4.4(b), TDD-RED-first per §11.4.43 / §11.4.115) BEFORE the pre-build sweep and main build proceed. The review MUST iterate (re-review after each remediation) until it surfaces no remaining blocking findings; only then does the build proceed. The review is itself anti-bluff: its conclusions are captured evidence (§11.4.5 / §11.4.69), and a review that rubber-stamps a defective batch is the same severity class as a PASS-bluff.

Multiple strong layers. The review is one of multiple strong layers of checks, complementing — never replacing — the pre-build gate sweep (§1 four-layer), the §11.4.92 multi-pass change-evaluation, the §11.4.108 four-layer fix-verification, the §11.4.110 pre-build build-readiness verdict + clash detector, and the post-build / runtime-on-clean-target / user-visible layers. §11.4.125 sits at the seam between "all batch work done" and "start the pre-build + build": it is the independent-review gate that the other layers do not provide (the multi-pass per §11.4.92 is self-review by the change author; §11.4.125 adds the structurally-separated reviewer seam per §11.4.70).

Composition. §11.4.125 composes with §11.4 (the end-user-quality covenant the review enforces), §11.4.1 (FAIL-bluffs equally caught), §11.4.4 (four-layer coverage of every remediation the review demands), §11.4.6 (no-guessing — review conclusions are FACT from captured evidence, never "looks fine"), §11.4.40 (the full-suite retest is a later gate; §11.4.125 is the pre-build review gate that precedes it), §11.4.43 (TDD-fix — remediation lands RED-first), §11.4.50 (deterministic-consistency of the tests the review validates), §11.4.70 / §11.4.20 (the review IS a subagent dispatch — structurally-separated reviewer eliminates self-review blind spots), §11.4.92 (multi-pass change-evaluation is the author-side pass; §11.4.125 is the independent-reviewer pass), §11.4.102 (any defect the review finds triggers systematic-debugging before remediation), §11.4.107 (the review verifies AV/test-validation analyzers are not themselves bluffs), §11.4.108 (the review checks each

fix's runtime signature is declared + verifiable), §11.4.110 (the build-readiness verdict + clash detector are mechanical gates the review complements with judgement).

Enforcement. Propagation gate `CM-COVENANT-114-125-PROPAGATION` enforces the literal anchor `11.4.125` across the consumer fleet (every `CLAUDE.md` / `AGENTS.md` / `QWEN.md`). Recommended gate `CM-CODE-REVIEW-GATE-BEFORE-BUILD` (the build orchestrator records a code-review-completed marker — with the review's findings + their remediation-or-no-blocking-findings determination — produced AFTER the batch's last fix and BEFORE the pre-build sweep + main build; a build that starts without a fresh review marker for the current batch FAILs). Paired §1.1 meta-test mutations: strip the `11.4.125` literal → propagation gate FAILs; start a build with no fresh code-review marker for the batch → `CM-CODE-REVIEW-GATE-BEFORE-BUILD` FAILs. (Gate-code implementation lands as a separate follow-up work item; this anchor defines the contract.)

Classification: universal (§11.4.17) — "an independent multi-layer code-review-agent gate analyzing all work + data + codebase + git history for quality/safety/will-it-really-work + genuine test-validation, with findings blocking the build, runs BEFORE the pre-build sweep + main build" is a platform-neutral discipline reusable by ANY project that builds an artifact from source; the consuming project supplies its concrete review-marker mechanism + build-orchestrator integration per §11.4.35.

Canonical authority: this Constitution.md §11.4.125 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--skip-code-review`, `--build-without-review`, `--no-review-gate`, `--review+optional`, `--trust-the-author` flag exists.

2026 06 04

§ . . — Default autonomous-loop
working mode from first prompt (User mandate,
— —)

Forensic anchor — verbatim user mandate (2026-06-04):

"Make sure that you continue work in endless fully autonomous loop, do not stop until new fully validated and verified version (tag) is created and published (all submodules and main repo) or IN A CASE OF some other main stream work until it is fully completed with all side work streams and nothing else is left in our working queue! THIS MUST BE ALWAYS the default working mode without us asking you! We tend to achieve ABSOLUTE EFFICIENCY, with this and all other projects which will incorporate this MANDATORY RULE / CONSTRAINT!!! This way of (your) working will be ALWAYS applied / followed / executed / fully respected, as soon as we assign / send first request (prompt) in the session! This stops only if we explicitly say so or nothing is left to be done in current working scope (release that will come / upcoming version)!!! Any mimicking (imitation) of this behavior / rules / mandatory constraints, false results or any kind of bluff(s) is ABSOLUTELY FORBIDDEN!!!!"

The endless fully-autonomous loop is the **DEFAULT working mode**, engaged automatically the moment the operator sends the FIRST request / prompt of a session — the operator MUST NOT have to ask for it, request it, restate it, or re-enable it per session. §11.4.87 framed the endless-loop covenant as an explicit-instruction opt-in ("continue in endless loop fully autonomously" or a semantically-equivalent phrasing); §11.4.126 is the **capstone** that promotes the same covenant to always-on: from the first prompt onward, every agent operates in the §11.4.87 loop discipline as the standing default, with §11.4.94 zero-idle, §11.4.97 maximum-idle-use, §11.4.101 autonomous-decision-over-blocking, and §11.4.103 continuous-parallel-stream all engaged by default — no per-session activation handshake.

The continuation contract (the loop continues until ONE of the two terminal conditions holds): (A) **Release scope** — a new, fully-validated-and-verified version (tag) is created AND published across all owned submodules AND the main repo to all configured remotes (per §2.1 multi-upstream push + §11.4.40 full-suite-retest-before-tag + §11.4.113 absolute-no-force-push merge-onto-latest-main). (B) **Non-release main-stream scope** — the main-stream work goal is fully completed AND every side work stream is done AND the working queue holds nothing left to

do for the current scope. Until ONE of (A) or (B) holds, the agent MUST keep working — it claims the next priority item, dispatches the next parallel stream, or progresses every non-blocked item, per §11.4.42 / §11.4.72 / §11.4.94 / §11.4.103.

The loop STOPS ONLY on: (1) the operator explicitly saying so (STOP / pause / end); (2) nothing left to do in the current working scope — the upcoming release / the current main-stream goal — with the queue genuinely empty per the (A)/(B) terminal conditions; (3) a §12 host-session-safety demand (the loop yields to host safety unconditionally). Idle-while-blocked parks one work unit, it does not stop the loop — the agent keeps progressing every non-blocked item in parallel per §11.4.101 + §11.4.94 + §11.4.97.

Goal — ABSOLUTE EFFICIENCY. The discipline exists to drive the project to completion with no operator-side restart overhead, no idle gaps, no stop-and-wait round-trips. It applies to this project AND every project that incorporates this Constitution (universal per §11.4.17).

Anti-bluff — actually do the work, never imitate it. Mimicking / imitating this loop behaviour, narrating continuation without performing it, fabricating progress, or emitting false / bluff results of ANY kind is ABSOLUTELY FORBIDDEN — this composes the entire §11.4 anti-bluff covenant family (§11.4 / §11.4.1 / §11.4.2 / §11.4.5 / §11.4.6 / §11.4.50 / §11.4.69 / §11.4.107). The agent MUST genuinely perform the continuous work and capture positive evidence for every closure; a report that claims the loop ran while no real work / no captured evidence was produced is a §11.4 PASS-bluff at the operating-mode layer, the severity class this mandate specifically prohibits.

Classification: universal (§11.4.17) — a platform-neutral operating-mode discipline reusable by ANY project; the consuming project supplies its concrete release-tag / publish mechanism + working-queue source per §11.4.35. Composes with §11.4.87 (the endless-loop covenant — §11.4.126 promotes it from opt-in to always-on default) / §11.4.94 (zero-idle priority-first parallel-by-default) / §11.4.97 (maximum-idle-use + progress-update cadence) / §11.4.101 (autonomous-decision-over-blocking) / §11.4.103 (continuous parallel-stream routine) / §11.4.66 (interactive clarification ONLY when a decision is genuinely operator-blocked per §11.4.101) / §11.4.6 (no-guessing — terminal conditions determined from captured queue/tracker state, never assumed) / §11.4.40 (release-tag full-suite retest is the (A) terminal gate) / §11.4.42 (iteration discipline is the loop body) / §11.4.72 (audio-first priority within the loop) / §11.4.113 (publish via merge-onto-latest-main, no force-

push) / §2.1 (multi-upstream publish) / §12 (host-session safety overrides the loop). Propagation gate CM-COVENANT-114-126-PROPAGATION (literal 11.4.126 across the consumer fleet) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: this Constitution.md §11.4.126 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no

```
11 4 127
--ask-before-continuing , --single-turn-only , --not-default-loop , --
mimic-OK flag exists.
2026 06 06
```

§ . . — Session-handoff resumption-
prompt mandate (User mandate,
— —)

Forensic anchor — verbatim user mandate (2026-06-06):

"make sure that in situations like this now when new session is needed you ALWAYS prepera such sentence - which will be valid for particular moment and the phase of the project and enough for work to continue."

When the agent determines a fresh session is needed (context-window limits, performance degradation) OR the operator asks whether a new session is needed / requests a handoff, the agent MUST ALWAYS prepare and provide a ready-to-paste **resumption prompt valid for that EXACT moment and project phase** — self-contained enough that pasting it into a fresh session resumes work with ZERO loss. The agent MUST proactively offer the resumption prompt the moment it recognises the need (not only when asked). Two variants on demand: a SHORT first-sentence variant ("Read <handoff docs> , then continue <terminal goal> ...") AND a FULL detailed block. The prompt MUST: (1) point to the live handoff doc(s) — .remember/remember.md if present + docs/CONTINUATION.md per §12.10 — and instruct reading them FIRST + git fetch --all ; (2) state the current PHASE + the immediate NEXT action + the terminal goal; (3) embed the exact current-state anchors (build IDs / artifact MD5, device/target serials, commit HEAD, in-flight PIDs + log paths, captured-evidence paths); (4) restate the binding constraints (anti-bluff §11.4 covenant, no-force-push §11.4.113, exact version/

naming, hardware/target gotchas); (5) be MOMENT-VALID — reflect the actual live state, NEVER a generic template. The handoff doc(s) MUST be brought current BEFORE the prompt is given (§12.10). A handoff that omits the resumption prompt, or gives a stale / generic one, is a §11.4.127 violation.

Classification: universal (§11.4.17) — a platform-neutral session-continuity discipline reusable by ANY project worked by context-bounded agents; the consuming project supplies its concrete handoff-doc paths + state anchors per §11.4.35. Composes with §12.10 (CONTINUATION doc — the resumption prompt is its operator-facing entry point) / §11.4.6 (no-guessing — the embedded state is captured fact, not assumed) / §11.4.66 (interactive clarification — "do you need a new session?" is itself a §11.4.66 decision point) / §11.4.87 / §11.4.103 / §11.4.126 (the fresh session CONTINUES the always-on autonomous loop — the resumption prompt is HOW the loop survives a context reset). Propagation gate CM-COVENANT-114-127-PROPAGATION (literal 11.4.127 across the consumer fleet) + recommended gate CM-HANDOFF-RESUMPTION-PROMPT-PRESENT + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: this Constitution.md §11.4.127 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no

^{11 4 128}
--skip-handoff-prompt , --generic-prompt-OK , --no-resumption-
^{2026 06 06}
sentence , --handoff-without-state flag exists.

§ . . — Always-on device-recording
mandate (User mandate, - -)

Forensic anchor — direct user mandate (2026-06-06): we MUST ALWAYS live-record all available data from all devices we use for testing or that we know manual testing is being performed on, done EXTRA carefully so the recording never harms the device, its performance, or causes side effects; recorded raw data MUST NOT be processed without need (token-conscious) and is ALWAYS excluded from version control + the code-intelligence index; only curated analysis/evidence is committed, and only during release preparation.

The mandate (ALL must hold). For EVERY test/debug device the project uses — and every device on which the operator is known to be performing manual testing — across EVERY reachable transport (USB or wireless ADB / SSH / serial / network introspection API), the project MUST ALWAYS live-record all available data sources it can later analyse: all activities, all logs (system log, kernel log, dropbox, crash/ANR traces), performance metrics (CPU / memory / I/O / thermal / load), every sink-side / downstream introspection report per §11.4.13, and any other live-changeable parameter. (1) **Extra-careful, side-effect-free.** Recording MUST be designed so it NEVER harms the device, degrades its performance, or causes a side effect — non-invasive read-only probes only, bounded sampling intervals, bounded write-volume, an observer-effect budget (no probe that mutates the system-under-test). A recorder that perturbs the device-under-test is a §11.4.128 violation, NOT evidence. (2) **Background + parallel + subagent-driven.** Recording MUST run in the background in parallel with the main work stream (per §11.4.103 continuous-parallel-stream), dispatched subagent-driven per §11.4.70 — never blocking the main stream, never the only thing the conductor does. (3) **Token-conscious — record-now, analyse-later.** Raw recorded data MUST NOT be processed / analysed without need; the default is capture-and-store. The ONLY standing exception is during release-tag preparation (the §11.4.40 full-suite retest / §11.4.42 release-gate) OR when the operator explicitly asks — at which point the relevant slice is analysed and the curated findings become captured-evidence. Analysing raw recordings continuously, eagerly, for no tracked reason burns tokens and is a §11.4.128 violation. (4) **Raw is git-ignored AND code-intelligence-excluded; only curated evidence is committed.** The raw recording corpus MUST be excluded from version control (with a §11.4.77 regeneration-mechanism declaration — the mechanism is "re-record by re-running the recorder against the device", with the honest caveat that a past device-state is not byte-reproducible) AND excluded from the code-intelligence index per §11.4.78/§11.4.79 (a raw log/recording corpus is not source; indexing it pollutes the graph). Only the CURATED analysis / extracted evidence artefacts are committed — and only during release preparation per §11.4.83 (docs/qa/<run-id>/) — never the raw corpus. (5) **Deterministic archive layout.** The raw corpus MUST be laid out by date → combined-source-state-hash → device → recording-sequence so any artefact is traceable to the exact code state and device it came from: <recording-root>/YYYY-MM-DD/<combined main+submodules state hash>/<DEVICE>_<SERIAL>/recording_NNN/<files> (monotonic NNN per device per state-hash). (6) **Anti-bluff.** A recorder claimed running but producing no growing corpus is a §11.4 bluff;

a curated finding cited as evidence MUST trace to a real raw-corpus path under the deterministic layout; the recorder's own health (alive + corpus growing) is itself a captured-evidence assertion per §11.4.5/§11.4.69.

Classification: universal (§11.4.17) — a platform-neutral test-evidence-capture discipline reusable by ANY project with reachable test/debug devices; the consuming project supplies its concrete device serials, transport probes, recorder helper, recording-root, and performance-metric set per §11.4.35. Composes with §11.4.2 (recorded-evidence — §11.4.128 guarantees the recording exists to BE the evidence) / §11.4.5 (captured-evidence quality) / §11.4.13 (sink-side introspection IS one recorded source) / §11.4.69 (sink-side positive-evidence taxonomy) / §11.4.40 + §11.4.42 (release-prep is the analyse-the-corpus trigger) / §11.4.70 (subagent-driven recorder dispatch) / §11.4.77 (raw corpus git-ignored WITH a regeneration-mechanism declaration) / §11.4.78 + §11.4.79 (raw corpus excluded from the code-intelligence index) / §11.4.83 (curated evidence lands under docs/qa/<run-id>/) / §11.4.103 (background parallel stream) / §11.4.119 (single-resource-owner — a read-only recorder never contends with an exclusive driver). Propagation gate CM-COVENANT-114-128-PROPAGATION (literal 11.4.128 across the consumer fleet) + recommended gate CM-DEVICE-RECORDING-ALWAYS-ON (recorder present + raw-corpus git-ignored + code-intelligence-excluded + deterministic layout) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: this Constitution.md §11.4.128 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--skip-recording` , `--record-without-layout` , `--commit-raw-corpus` , `--index-raw-corpus` , `--analyse-corpus-always` , `--invasive-probe-OK` flag exists.

§ . . — Huge-blocker release protocol (User mandate, - -)

Forensic anchor — direct user mandate (2026-06-06): when a huge blocker is discovered during release validation we MUST stop all testing, return to fixing ALL discovered issues, process all recorded data from the last session, land rock-solid fixes, author NEW validation+verification tests of ALL supported test types, rebuild,

reflash, and RESTART the full validation+verification of every fix/change from the last release tag to now — on both devices in parallel, recorded, with real physical captured proofs and no bluff.

When a HUGE BLOCKER (a release-blocking-severity defect — a core user-facing capability broken, a regression that invalidates the in-flight validation cycle, or a defect whose blast radius reaches the batch's other fixes) is discovered during release validation, the agent **MUST** execute the following protocol in order, with NO partial / spot-check shortcut: (1) **STOP all testing** immediately on every device (the §11.4.4 test-interrupt-on-discovery STOP, escalated to release granularity — a validation cycle that continues past a huge blocker produces "all green" summaries on a known-broken build, the exact §11.4 PASS-bluff). (2) **Return to fixing ALL discovered issues** — not only the huge blocker but every defect surfaced this cycle; investigate each to root cause per §11.4.102 systematic-debugging, isolate regressions against the last known-good tag per §11.4.114. (3) **Process all recorded data from the last session** — analyse the §11.4.128 raw-corpus slice for the cycle (this is the §11.4.128(3) release-prep analyse-trigger) so no defect the recordings already captured is missed. (4) **Land rock-solid fixes** — per §11.4.123 (rock-solid proof or deep research), §11.4.43/§11.4.115 (RED-on-broken-artifact then GREEN), §11.4.9 batch-source-fixes-before-rebuild. (5) **Author NEW validation+verification tests of ALL supported test types** — per §11.4.27 (unit / integration / e2e / full-automation / security / stress / chaos / performance / UI / UX / Challenges / HelixQA) + §11.4.85 stress+chaos, each anti-bluff with captured evidence + paired §1.1 mutation. (6) **Rebuild + reflash** — full rebuild (NOT module-only) for retest-baseline integrity, deploy to a CLEAN target per §11.4.108 (no stale-overlay shadow). (7) **RESTART the full validation+verification from the last release tag to now** — re-validate EVERY fix/change landed since the last tag (per §11.4.40 full-suite-retest-before-tag — the huge-blocker discovery invalidated the in-flight cycle, so the cycle restarts from the beginning, NOT resumes), running on both/all owned devices IN PARALLEL per §11.4.103/§11.4.119 (single-resource-owner partitioning), every device's run RECORDED per §11.4.128, producing REAL physical captured proofs per §11.4.5/§11.4.69/§11.4.107 with no bluff of any kind. This anchor BINDS the existing release-cycle anchors for the huge-blocker case — it adds the STOP → fix-all → process-recordings → new-tests-all-types → rebuild → reflash → full-restart sequencing and the "restart, never resume" rule, citing §11.4.4 + §11.4.40 + §11.4.108 + §11.4.114/§11.4.115 + §11.4.123 + §11.4.128 rather than duplicating their text.

Classification: universal (§11.4.17) — a platform-neutral release-discipline reusable by ANY project that runs a release-validation cycle; the consuming project supplies its concrete device set, build/flash mechanism, and last-known-good tag per §11.4.35. Composes with §11.4.4 (test-interrupt-on-discovery — §11.4.129 is its release-granularity escalation) / §11.4.40 (full-suite retest before tag — §11.4.129 mandates the RESTART, not resume) / §11.4.42 (iteration discipline) / §11.4.9 (batch-source-fixes-before-rebuild) / §11.4.27 (all-test-types coverage) / §11.4.85 (stress+chaos) / §11.4.102 (systematic-debugging) / §11.4.108 (clean-target rebuild+reflush+runtime-signature) / §11.4.114 + §11.4.115 (regression isolation + RED-on-broken-artifact) / §11.4.123 (rock-solid proof) / §11.4.128 (process the recorded corpus) / §11.4.103 + §11.4.119 (parallel recorded multi-device).

Propagation gate `CM-COVENANT-114-129-PROPAGATION` (literal `11.4.129` across the consumer fleet) + recommended gate `CM-HUGE-BLOCKER-FULL-RESTART` (a release-blocker discovery during validation triggers a from-the-last-tag full restart, not a resume) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: this Constitution.md §11.4.129 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no

```
--resume-after-blocker , --spot-validate-after-fix ,  
11 4 130  
--skip-recording-analysis , --skip-new-tests , --module-only-after-  
blocker , --single-device-restart flag exists.  
2026 06 06
```

§ . . — Post-remediation validate-the-fix-FIRST-after-redeploy (User mandate, - -)

Forensic anchor — direct user mandate (2026-06-06): when a blocker discovered during release validation is fixed and a new artifact is produced + redeployed, we MUST first re-test the SPECIFIC last-failing features and validate+verify the just-incorporated fixes BEFORE proceeding to the broader / full validation — a first fix attempt may not work, may be incomplete, or may regress again, so confirming the targeted fix on the new artifact first prevents wasting a full multi-hour cycle on a still-broken fix.

When a blocker / critical failure is discovered during release validation and then FIXED, and a new artifact is produced (rebuild / new flashing image / redeploy) and the target is reflashed / redistributed / updated, the agent MUST execute the following ordering with NO shortcut: (1) **Re-test the SPECIFIC last-failing features FIRST** — run the targeted guard tests for exactly the defects this fix addressed, BEFORE any broader / full-suite validation. (2) **Validate + verify the just-incorporated fixes with real captured evidence** — the RED test that reproduced the defect on the pre-fix artifact (per §11.4.115) MUST now flip GREEN at `RED_MODE=0` on the new artifact, AND the fix's declared runtime signature (per §11.4.108) MUST verify on the CLEAN target the redeploy produced. Metadata-only / config-only / absence-of-error / grep-without-runtime PASS are forbidden (§11.4 / §11.4.1) — the proof is captured per §11.4.5 / §11.4.69 / §11.4.107 / §11.4.123. (3) **Only after the targeted fix is CONFIRMED working on the new artifact** does the agent proceed to the §11.4.40 full retest of every fix/change from the last release tag to now.

Rationale. A first fix attempt may not work, may be incomplete, or may regress again under the new artifact. Confirming the targeted fix FIRST catches a fix-did-not-take case immediately — instead of discovering it hours later at the END of a full multi-hour validation cycle (and then having to restart the whole cycle per §11.4.129). The targeted re-test is cheap; the full cycle is expensive; ordering cheap-confirmation-first is the §11.4.82 iteration-speedup discipline applied to the post-blocker reflash. This is the §11.4.46 recent-work-validation gate specialised for the post-blocker-reflash case: the §11.4.129 huge-blocker protocol's step (7) full-restart MUST be preceded by this targeted-fix-first confirmation phase.

Honest boundary (§11.4.6): "the fix probably took" is not "the fix took" — the targeted RED → GREEN flip + runtime-signature verification on the new artifact is the captured proof; proceeding to the full cycle without it is a §11.4.130 violation. A targeted re-test that still FAILs on the new artifact is a finding — return to the §11.4.114/§11.4.115 isolate → RED → fix loop and produce another artifact; do NOT proceed to the full cycle on a still-broken fix.

Classification: universal (§11.4.17) — a platform-neutral release-discipline reusable by ANY project that fixes a blocker, redeploys, and re-validates; the consuming project supplies its concrete targeted-guard-test set, build/flash mechanism, and clean-deployment / runtime-signature observables per §11.4.35. Composes with §11.4.4 (test-interrupt-on-discovery — the blocker that triggered the fix) / §11.4.40

(full-suite retest before tag — §11.4.130 is the targeted-first phase that GATES the full retest) / §11.4.46 (validate-recent-work-before-post-flash-tests — §11.4.130 is its post-blocker-reflash specialisation) / §11.4.108 (runtime-signature-as-definition-of-done on the clean target the redeploy produced) / §11.4.114 (regression isolation — a still-failing targeted re-test re-enters the isolation loop) / §11.4.115 (RED-baseline + polarity-switch — the RED test flips GREEN on the new artifact) / §11.4.123 (rock-solid proof — captured evidence, never "probably took") / §11.4.129 (huge-blocker protocol — §11.4.130 is its post-reflash targeted-confirmation phase that precedes the step-7 full-restart) / §11.4.82 (iteration-speedup — cheap-confirmation-first). Propagation gate CM-COVENANT-114-130-PROPAGATION (literal 11.4.130 across the consumer fleet) + recommended gate CM-FIX-FIRST-AFTER-REDEPLOY (after a post-blocker redeploy, the targeted last-failing-feature guard tests run + pass on the new artifact before the full retest starts) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: this Constitution.md §11.4.130 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--skip-targeted-retest`, `--full-cycle-first`, `--assume-fix-took`, `--validate-fix-at-end`, `--skip-red-green-flip-on-new-artifact` flag exists.

§ . . — Standing session-resumption
file mandate (User mandate,
— —)

Forensic anchor — verbatim user mandate (2026-06-07): "Make this markdown a standard file which will be written EVERY TIME when we need fresh session out of the box! It MUST BE always up to date and in sync so whenever new session is created all we have to do is just point to it!"

Every project MUST maintain a SINGLE canonical, always-current **session-resumption file** at a fixed, project-declared standard path (the project declares the exact path once per §11.4.35 and never moves it without a §11.4.66 operator decision). This file is the OUT-OF-THE-BOX entry point for any fresh session: creating a new session requires ONLY pointing the new agent at this one file. §11.4.131 promotes §11.4.127 (which mandated PREPARING a resumption prompt

on demand) into a **STANDING**, version-controlled **ARTIFACT** — the resumption prompt is no longer an ephemeral in-message reply but a persistent file that is **ALWAYS** present and **ALWAYS** in sync.

(A) **Existence + fixed path.** The file **MUST** exist at the declared canonical path at all times; its location is encoded as a literal path string in the project-layer instantiation (§11.4.35) and never silently moved.

(B) **Always written + always synced.** The file **MUST** be (re)written whenever a fresh session is needed **OR** whenever the live project state materially changes (new **HEAD**, build/artifact id, phase transition, device/target state, in-flight job, blocking decision) — the §12.10 **CONTINUATION.md** trigger set applied to the resumption file. A stale resumption file (describing a prior **HEAD** / phase / artifact) is a §11.4.131 violation of the same severity class as a §12.10 stale-**CONTINUATION** violation: it silently breaks the zero-loss-resumption guarantee.

(C) **Content (composes §11.4.127).** The file **MUST** carry both the §11.4.127 **SHORT** (one-paste first-sentence) **AND FULL** (paste-ready block) variants; point to the live handoff docs read **FIRST** (`.remember/remember.md` if present + `docs/CONTINUATION.md` per §12.10) + `git fetch` ; embed the exact live-state anchors (**HEAD** commit, build/artifact ids + checksums, device/target serials, in-flight PIDs + log paths, captured-evidence paths); state the current **PHASE** + immediate **NEXT** action + terminal goal; and restate the binding constraints (anti-bluff §11.4, no-force-push §11.4.113, exact version/naming, hardware/target gotchas). It **MUST** be **MOMENT-VALID**, never a generic template (§11.4.6).

(D) **Export + freshness.** The file is in §11.4.65 scope → synchronized `.html` / `.pdf` siblings refreshed on every update; a §11.4.44 revision header tracks its monotonic revision + last-modified timestamp.

(E) **Out-of-the-box resumption.** A fresh session, given **ONLY** this file's path, **MUST** be able to fully resume with zero additional context. The file **IS** the session bootstrap.

Composes §12.10 (**CONTINUATION.md** live-state — §11.4.131 is the dedicated paste-ready-prompt sibling of it) / §11.4.127 (resumption-prompt discipline promoted to a standing file) / §11.4.65 (universal export) / §11.4.44 (revision header) / §11.4.6 (moment-valid, never generic) / §11.4.66 (path change is an operator decision) / §11.4.126 (default-loop continuity across sessions).

Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-131-`

PROPAGATION (literal 11.4.131 across the consumer fleet) + recommended gate CM-SESSION-RESUMPTION-FILE-PRESENT (the declared file exists + carries SHORT+FULL + revision header + is fresh vs HEAD) + paired §1.1 meta-test mutation (strip the file / make it stale → gate FAILs; gate-code = separate work item).

Canonical authority: this Constitution.md §11.4.131 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no

```
11 4 132
--skip-resumption-file , --ephemeral-prompt-only ,
--stale-resumption-OK , --generic-template-OK flag exists.
2026 06 07
```

§ . . — Risk-ordered validation priority mandate (User mandate, - -)

Forensic anchor — verbatim user mandate (2026-06-07): "We MUST ALWAYS first test and validate features, functionalities and fixes/changes that have been worked most recently, the ones which were most problematic, which have the most chance to crash or break again, the ones which have been re-opened the most times! Then, after we validate and verify all this with real (physical) proofs and hard evidence, with no false results and bluffs of any kind, we continue with all other existing tests in the test suites! This IS MANDATORY."

Tests, validations, and verifications MUST run in **RISK-DESCENDING order** — the highest-risk set FIRST, and only AFTER that set is fully GREEN with real (physical) captured evidence does the remainder of the suite run. The risk ranking is computed from a closed set of factors, highest-risk first:

- (a) **Most-recently-worked** — features / functionalities / fixes / changes touched most recently (the freshest code is the least-exercised in the wild and the most likely to carry an undiscovered defect).
- (b) **Historically most-problematic** — the items with the longest defect history, the most prior fixes, the most prior failures.
- (c) **Highest crash/break/regress likelihood** — items whose blast radius, complexity, or dependency surface gives them the greatest chance to crash or break again.

(d) **Most-reopened** — items with the highest reopen count per §11.4.55 (a high reopens-count is the strongest empirical signal of fragility — a defect that keeps coming back).

The highest-risk set, ordered by (a)–(d), is validated FIRST. Each item in that set MUST pass with real (physical) captured evidence per §11.4.5 / §11.4.69 / §11.4.107 — no metadata-only / config-only / absence-of-error / grep-without-runtime PASS (§11.4 / §11.4.1), no false results, no bluff of any kind (§11.4.6). ONLY AFTER the entire highest-risk set is GREEN with captured proof does the agent proceed to all other existing tests in the test suites. Running the full suite in arbitrary order, or running lower-risk tests before the highest-risk set is GREEN, is a §11.4.132 violation: it spends the operator's scarce validation window on the items least likely to fail while deferring the items most likely to fail.

§11.4.132 REFINES / STRENGTHENS §11.4.130 (post-remediation validate-the-fix-FIRST-after-redeploy — §11.4.130 says re-test the just-fixed items first; §11.4.132 generalises "first" to the full risk-ordered set) + §11.4.46 (validate-recent-work-before-post-flash-tests — §11.4.46 gates the full suite on recent-work GREEN; §11.4.132 adds the explicit risk-ordering WITHIN the recent / high-risk set) + §11.4.42 (iteration-discipline priority order — §11.4.132 applies the same priority discipline to VALIDATION ordering, not just implementation ordering). The risk ranking is the validation-layer analogue of the §11.4.42 / §11.4.72 implementation-layer priority queue.

Classification: universal (§11.4.17) — a platform-neutral validation-ordering discipline reusable by ANY project; the consuming project supplies its concrete recency / problematic-history / reopen-count sources (e.g. the §11.4.93 workable-items DB `reopens_count` + `last_modified` columns) per §11.4.35. Composes §11.4.4 / §11.4.5 / §11.4.6 / §11.4.7 / §11.4.40 / §11.4.42 / §11.4.46 / §11.4.50 / §11.4.55 / §11.4.69 / §11.4.107 / §11.4.130. Propagation gate `CM-COVENANT-114-132-PROPAGATION` (literal `11.4.132`) + recommended gate `CM-RISK-ORDERED-VALIDATION-PRIORITY` + paired §1.1 meta-test mutation (gate-code = separate work item).

Canonical authority: this Constitution.md §11.4.132 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--skip-risk-ordering`, `--any-order-OK`, `--suite-order-fixed` flag exists.

§ . . – Target-System + hardware safety mandate (User mandate, – –)

Forensic anchor — verbatim user mandate (2026-06-08): "Make sure that all changes we do to the System are ALWAYS safe for the System itself and for the hardware the system runs on! This is MANDATORY."

Every change to the TARGET system — firmware, kernel, init / boot scripts, drivers, sysfs / devfreq / voltage / clock / thermal / regulator register writes, partition layout / bootloader / U-Boot, HAL, framework, prebuilts, device configuration, or any equivalent on a non-Android target — MUST ALWAYS be safe for BOTH: (a) **the target System itself** — the change MUST NOT brick the device, induce a boot-loop, corrupt user / system data, or render the device unrecoverable; AND (b) **the hardware the System runs on** — the change MUST NOT drive any hardware-control surface beyond its safe electrical / thermal / voltage / clock envelope, and MUST NOT damage panels, storage, radios, regulators, or any other physical component.

Concrete obligations (ALL hold):

1. **Reversible-first.** Prefer reversible changes; when a change is irreversible or high-blast-radius (bootloader / U-Boot / partition layout / one-way fuse / OTP), verify it against a known-good value (MD5 / checksum / golden manifest) BEFORE applying, and capture a pre-op backup of any state the change overwrites (per §9.2). A wrong irreversible write that bricks the device is the worst-case the mandate exists to prevent.
2. **No unverified hardware-control writes.** NEVER write an unverified value to a hardware-control sysfs node / register / firmware field that governs voltage, clock frequency, regulator output, thermal-throttle threshold, current limit, or any other physical-safety parameter. Every such value MUST be within the component's datasheet / vendor-documented safe range; the safe range MUST be established as FACT (cited datasheet / vendor reference per §11.4.8 + §11.4.99), never guessed (§11.4.6). Pinning a max OPP / disabling a throttle / raising a regulator output without proving the cooling design and electrical envelope tolerate it is forbidden.

3. **Thermal / performance changes respect the cooling design.** Any change that raises sustained power, frequency, or voltage (e.g. forcing a performance governor, pinning the top OPP, disabling idle / thermal management) **MUST** be validated against the device's actual cooling capacity — captured thermal evidence (sustained-load temperature $\leq$ the documented limit) is the proof; "it boots" is not proof it is thermally safe.
4. **Flashing uses the sanctioned tool + verified image.** Deploying firmware **MUST** use the project's sanctioned flashing tool against a freshly-built, integrity-verified image — never an ad-hoc partition write, never a stale / unverified artifact, never a tool / image whose provenance is unestablished.
5. **Unprovable-safety $\Rightarrow$ blocked.** A change whose safety for the target System and its hardware cannot be established from captured evidence **MUST** be treated as UNSAFE and blocked — proceed only after the safe outcome is determined as FACT (§11.4.6), the change is reversible OR backed up (§11.4.101 reversible-first + §9.2), and (when the safe choice still cannot be determined) the operator is asked via §11.4.66 / §11.4.101's block-only-when rule. "Probably safe" is the exact guess this mandate forbids.

Distinct from §12 host-session safety. §12 (and §12.6 / §12.7 / §12.8 / §12.9 / §12.10) protects the DEVELOPER's HOST machine and session (no host suspend / logout / OOM cascade / unbounded-memory build). §11.4.133 protects the TARGET device the System is built FOR and runs ON. Both apply; neither weakens the other. A build step can be host-safe (bounded memory, no host suspend) yet still ship a target-unsafe change (an over-voltage regulator write) — §11.4.133 is the gate for the latter.

Classification: universal (§11.4.17) — a platform-neutral target-safety discipline reusable by ANY project that produces firmware / kernel / driver / device-config changes deployed onto target hardware; the consuming project supplies its concrete hardware-control surfaces, datasheet-safe ranges, known-good bootloader / image hashes, and sanctioned flashing tool per §11.4.35. Composes §12 (host-session safety — sibling-not-overlapping protection scope) / §11.4.6 (no-guessing — the safe range / known-good value is FACT, never assumed) / §11.4.101 (reversible-first + autonomous-decision-over-blocking — irreversible high-blast-radius hardware writes are the class that escalates to an operator question) / §11.4.108 (the change's runtime signature on a clean target proves it landed AND is safe) / §11.4.123 (rock-solid proof — safety claims need captured

evidence, never a metadata-only assertion). Propagation gate `CM-COVENANT-114-133-PROPAGATION` (literal `11.4.133` across the consumer fleet) + recommended gate `CM-TARGET-HARDWARE-SAFETY` + paired §1.1 meta-test mutation (gate-code = separate work item).

Canonical authority: this Constitution.md §11.4.133 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no

`--unsafe-hardware-write`, `--skip-system-safety`, `--brick-risk-accepted` flag exists.

2026 06 08

§ . . — Code-review iterate-until-GO
+ rock-solid-evidence mandate (User mandate,
- -)

Forensic anchor — verbatim user mandate (2026-06-08): "For any fixes/changes given back to us for re-work by the code-review process, once we fix/improve everything per the code-review's requests, we MUST RE-RUN code-review AGAIN until we get a GO from it with NO new issues reported or warnings of any kind! All results produced by this whole process MUST ALWAYS give us rock-solid PHYSICAL evidence that the fixed/improved codebase really works now as expected, with no false results and no bluff(s) of any kind."

When the §11.4.125 code-review returns ANY finding — BLOCKING, nit, or warning — and the author fixes/improves the batch per that review, the code review MUST BE RE-RUN, and MUST KEEP being re-run after each remediation round, until it returns a clean GO with ZERO new issues AND ZERO warnings of any kind. A single pass that "addressed the findings" is NOT sufficient: the corrected batch MUST pass a FRESH adversarial review (a re-review can surface NEW findings introduced by the very fixes that closed the prior ones — the §11.4.1 fix-A-creates-B failure mode). The loop terminates ONLY on a clean GO (no new findings, no warnings); a residual warning is itself a finding that re-arms the loop.

Every round's verdict AND every fix's validation MUST carry rock-solid PHYSICAL captured evidence per §11.4.5 / §11.4.69 / §11.4.107 (captured audio / video / sysfs / dumphsys / sink-side / runtime-signature) proving the fixed/improved codebase REALLY works as expected on the target — never metadata-only /

configuration-only / absence-of-error / grep-without-runtime. No false results, no bluff of any kind, at any round. A reported GO unbacked by captured physical evidence is itself a §11.4 PASS-bluff at the review-loop layer.

§11.4.134 REFINES / STRENGTHENS §11.4.125 (which requires the review to iterate "until no blocking findings remain"): it makes the loop EXPLICIT (re-run after every remediation round, not once), raises the termination bar to ZERO findings AND ZERO warnings (not merely zero-blocking), and BINDS rock-solid physical evidence to every round (not just the final one). Classification: universal (§11.4.17). Composes §11.4.125 / §11.4.1 / §11.4.4 / §11.4.5 / §11.4.6 / §11.4.69 / §11.4.107 / §11.4.50 / §11.4.108 / §11.4.123. Propagation gate CM-COVENANT-114-134-PROPAGATION (literal 11.4.134) + recommended gate CM-CODE-REVIEW-ITERATE-UNTIL-GO + paired §1.1 meta-test mutation (gate-code = separate work item).

Canonical authority: this Constitution.md §11.4.134 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--skip-rereview`, `--single-review-pass`, `--warnings-ok`, `--evidence-optional` flag exists.

2026 06 08

§ . . — Standing regression-guard suite + every-fixed-defect-gets-a-permanent-regression-test (User mandate, - -)

Forensic anchor (the canonical case this anchor exists to prevent, FACT): the wrong-subtitle-on-2nd-display defect was "fixed" in a prior release via a source-side `CONTROL_MENU_LABEL_DENYLIST` that NO test mirrored or re-ran, so the NEXT chrome class (a settings-menu label) recurred silently while the GREEN suite passed.

Every project MUST maintain a STANDING regression-guard suite that runs on EVERY build+deploy and BLOCKS the release tag on any failure. Every closed defect (the project's stable ticket id, e.g. ATM-NNN) MUST, in the SAME commit as its fix (extending the §11.4.43 DOCUMENT step), register a permanent §11.4.115 RED-on-broken-artifact regression test into the suite — `RED_MODE=1` capturing the historical defect on a pre-fix artifact (the proof the guard is real), `RED_MODE=0` the

standing GREEN guard asserting the defect is ABSENT. A closure without a registered guard is a §11.4.123 violation (closure without permanent regression proof). The suite runs FIRST in the post-deploy cycle (highest-risk set per §11.4.132 risk-ordering) and is a §11.4.40 release-gate blocker. This is industry-standard bug-driven testing (cite: Google for Developers content-driven testing <https://developers.google.com/solutions/content-driven/backend/testing>; AOSP CTS/Tradefed continuous-regression model <https://source.android.com/docs/compatibility/cts>) made mechanical + enforced.

Classification: universal (§11.4.17). Composes §11.4.4 / §11.4.40 / §11.4.43 / §11.4.46 / §11.4.50 / §11.4.107 / §11.4.108 / §11.4.115 / §11.4.118 / §11.4.123 / §11.4.124 / §11.4.130 / §11.4.132. Recommended gates `CM-REGRESSION-GUARD-REGISTERED` + `CM-REGRESSION-GUARD-SUITE-WIRED` + `CM-COVENANT-114-135-PROPAGATION` (literal `11.4.135`) + paired §1.1 meta-test mutations (gate-code = separate work item).

Verdict-coverage at the release seam — ABSENCE of a verdict blocks exactly as a FAIL does (extension, research-derived, 2026-07-17). Forensic FACT (genericised, 2026-07-17): a consuming project's standing guard suite computed its exit code from ONLY its FAIL count — a guard that was never run (fix never built/ deployed → `REQUIRES-REBUILD SKIP` → `PENDING`) affected nothing; measured live on a release-candidate target, 61 registered guards produced `PASS 0 / FAIL 5 / PENDING 56` and the suite exited `0 "not blocked"`, and the release-tag tooling consulted the guard suite `ZERO` times — to the release gate, "never verified" and "verified good" were the same colour. That single fail-open explains, without needing hundreds of separate bugs, how 95% of that tracker's done-claiming items shipped unguarded and how the recorded reopen-per-recorded-fix rate reached 52%. Therefore (ALL hold): (1) **the suite's exit semantics MUST incorporate COVERAGE, never only the FAIL count** — a `PENDING/never-run` verdict is not a pass AND not ignorable at the seams that ship; (2) **the release seam blocks on verdict coverage of the release-candidate artifact** — the tag is refused unless, for the candidate's artifact fingerprint, every registered guard whose required topology is present on at least one release-validation target has produced a real `PASS-or-FAIL` verdict, and zero FAILs: `uncovered = registered ^ topology-present ^ no-verdict-for-this-fingerprint` ; refuse while `uncovered ≠ ∅ ∨ N_FAIL > 0` ; (3) **the release-tag tooling MUST consult the suite/verdict store mechanically** — a release path that never reads the guards is itself a §11.4 bluff surface regardless of how green the suite runs elsewhere; (4) **the**

two legitimate PENDING classes are handled without a fail-open AND without a fail-closed wall: on intermediate artifacts a REQUIRES-REBUILD PENDING is honest and non-blocking, but on the release candidate it is a contradiction — either the wrong artifact is deployed (re-run on the candidate) or the fix genuinely missed the build (then its item MUST NOT be in the release's done-set); both resolutions are forced, neither is a permanent exemption. Topology-absent guards are exempt from coverage ONLY via a checked-in target-pool topology map and MUST be enumerated in the release changelog as honest coverage gaps (§11.4.3 / §11.4.118) — never a silent exemption; a naive "PENDING always blocks" rule is REJECTED as a §11.4.201 FAIL-bluff (it would block every honest intermediate run and teach operators to disable the gate); (5) **adoption ratchet** — a pre-existing unguarded done-claiming backlog is snapshotted ONCE into a committed ratchet file and may only shrink (the live unguarded set must be a subset of the snapshot AND the snapshot may not gain lines vs its committed version), so day one is green, the disease cannot spread, and every reopened/re-touched item exits the ratchet permanently. Recommended gate `CM-RELEASE-VERDICT-COVERAGE` — self-tested per §11.4.205 / §11.4.107(10): golden-bad fixture = a topology-present guard with no verdict for the candidate fingerprint → the release gate MUST FAIL; golden-good = a complete verdict set → MUST PASS; negative control = a properly-enumerated topology-absent guard → MUST PASS (the false-positive guard) — + paired §1.1 mutation (revert the exit semantics to FAIL-count-only, or de-wire the release-tag consultation → the gate FAILs; gate-code = separate work item). Composes §11.4.3 / §11.4.40 / §11.4.107(10) / §11.4.115(F) / §11.4.118 / §11.4.146(D3) / §11.4.201 / §11.4.205.

Canonical authority: constitution submodule `Constitution.md` §11.4.135. Non-compliance is a release blocker. No escape hatch — no `--skip-regression-guard`, `--no-guard-on-close`, `--guard-optional` flag — and, per the 2026-07-17 extension, no `--pending-does-not-block-the-candidate`, `--fail-count-only-exit`, `--release-without-consulting-guards`, `--silent-topology-exemption` flag.

§ 11.4.136 – Real-content end-to-end playback-test mandate (User mandate, – –)

Refines/strengthens §11.4.107. Any test asserting media playback works MUST drive REAL content (catalog stream or offline reference clip) through the user's path (§11.4.48 UI-driven → §11.4.117 CV/OCR fallback) and assert it genuinely PLAYS via the §11.4.107 liveness battery PLUS a decoder-health census — a numeric drop-buffer budget, no buffer-timestamp re-order/discard, no codec-reject (cite: Android/Media3 ExoPlayer OEM pre-OTA playback-test mandate <https://developer.android.com/media/media3/exoplayer/oems> — "too many dropped buffers" >25, "unexpected presentation timestamp", "test timed out" failure modes). Metadata-only / launch-only / registration-only / single-frame / config-only PASS is forbidden (§11.4 / §11.4.1). A golden/reference clip corpus (cite: BBC ExoPlayer testing samples <https://github.com/bbc/exoplayer-testing-samples>; USPTO 8,387,015 empirical media-player testing) is the offline ground-truth.

Classification: universal (§11.4.17). Composes §11.4.5 / §11.4.48 / §11.4.50 / §11.4.107 / §11.4.117 / §11.4.123 / §11.4.13 / §11.4.69. Recommended gate CM-REAL-CONTENT-PLAYBACK-TEST + CM-COVENANT-114-136-PROPAGATION (literal 11.4.136) + paired §1.1 meta-test mutation (gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md §11.4.136. Non-compliance is a release blocker. No escape hatch — no --launch-proves-playback^{11 4 137} , --skip-decoder-health , --metadata-playback-pass-suffices flag.

§ 11.4.137 – Subtitle/caption content-correctness oracle + secure-display-proxy-honesty mandate (User mandate, – –)

Refines §11.4.117 (pixel-oracle) + §11.4.107 (liveness) + §11.4.112 (structural-impossibility).

Forensic anchor (FACT): tests tasked to "physically verify the 2nd-display (or 1st where no TV) subtitle" PASSEd GREEN while subtitles did NOT show / showed WRONG, because (a) the streaming app's player surface is FLAG_SECURE so screencap -d <secondary> returns BLACK (cite: Android Developers <https://>

developer.android.com/security/fraud-prevention/activities; Nightwatch <https://www.nightwatchcybersecurity.com/2016/04/13/research-securing-android-applications-from-screen-capture/>; WithSecure <https://labs.withsecure.com/advisories/screencapture-via-ui-overlays-in-mediaprojection>) → autonomous PIXEL verification is structurally impossible (§11.4.112), so the test fell back to a non-pixel proxy (the accessibility-scraped forwarded cue / `persist.atmosphere.subdebug logcat`), and (b) the proxy's validation accepted a chrome/menu LABEL (`Аудио и субтитры`) as a valid subtitle because the prose floor accepted any multibyte prose and NO menu-label denylist + NO position/cadence check existed.

The mandate: a test asserting a subtitle/caption is correct MUST classify the *content class* of the rendered/forwarded cue — a present cue is NOT a correct cue. A cue is CHROME (FAIL) if it is a known control/menu label (a closed multilingual deny-list MIRRORED from the source-side denylist, case-folded both sides incl. non-ASCII scripts), is time-counter/numeric chrome, is NOT prose, is OUTSIDE the lower safe-title position band where caption position is available (cite CEA-708 9-anchor caption-window grid: captions live in the lower-center safe band, chrome elsewhere <https://www.gte-media.com/cea-608-and-cea-708-captions/>), OR is STATIC across the window (a real subtitle changes → dialogue cadence; a menu label does not → ≥ 2 distinct prose cues required, a metamorphic relation per https://en.wikipedia.org/wiki/Metamorphic_testing). A cue is DIALOGUE (PASS) only when prose, not-denied, not-chrome, position-ok, AND cadence ≥ 2 OR it fuzzy-matches the SOURCE-extracted expected cue via normalized edit distance (§11.4.123 host-side ground truth; cite SubER <https://github.com/apptek/SubER>). The content oracle MUST be self-validated golden-good/golden-bad (§11.4.107(10)) and the test-side deny-list MUST be verified present in the SHIPPED artifact (§11.4.108) — a source-green denylist with no test mirror + no artifact check is the exact recurrence pattern this anchor forbids.

Secure-display-proxy honesty (§11.4.112): where FLAG_SECURE makes pixel verification structurally impossible, the rock-solid autonomous proof is the player's own caption telemetry (the subdebug/Cue ground-truth channel) + source-track presence + content-class oracle — NEVER a faked pixel "physical" pass; the human-eye pixel confirmation is `operator_attended` (§11.4.52) with a tracked migration item. App-agnostic (keys off content class, not per-app rule) so it is the standing guard for the whole roster.

Classification: universal (§11.4.17). Composes §11.4.3 / §11.4.5 / §11.4.6 / §11.4.107 / §11.4.108 / §11.4.112 / §11.4.115 / §11.4.117 / §11.4.123 / §11.4.13 / §11.4.69. Recommended gate `CM-SUBTITLE-CONTENT-CORRECTNESS-ORACLE` + `CM-COVENANT-114-137-PROPAGATION` (literal `11.4.137`) + paired §1.1 meta-test mutation (strip the denylist/position/cadence check → golden-bad `Аудио и субтитры` `PASSes` → gate FAILs).

Canonical authority: constitution submodule `Constitution.md` §11.4.137. Non-compliance is a release blocker. No escape hatch — no `--present-cue-is-correct`, `--skip-chrome-oracle`, `--length-heuristic-suffices`, `--pixel-pass-on-secure-display`, `--skip-position-check`, `--skip-cadence-check` flag.
11 4 138
2026 06 08

§ . . — Operator-escape => mandatory bluff-audit + permanent guard (User mandate, — —)

When the operator (or any out-of-band channel) finds a defect that the GREEN test suite passed, this is by definition a §11.4 PASS-bluff — it MUST trigger, before the fix is closed: (1) a §11.4.102 systematic-debugging pass to FACT-root-cause; (2) a bluff-audit identifying the EXACT assertion that should have caught it but didn't, cited to `file:line` (canonical example: `lib/subtitle_content_validation.sh:sub_is_prose()` returning TRUE for `Аудио и субтитры`); (3) a permanent §11.4.135 regression guard registered in the SAME commit as the fix, with its §11.4.115 RED capturing the operator-found defect; (4) the bluff-audit committed under `docs/research/<scope>/<defect>_bluff_audit/`. Closing an operator-found defect WITHOUT the bluff-audit + permanent guard is itself a §11.4 violation (the bluff that let it through is still live and the defect will recur — the operator's "repeatedly reworking + reopening the same issues with no real progress" pain).

Classification: universal (§11.4.17). Composes §11.4 / §11.4.1 / §11.4.102 / §11.4.108 / §11.4.115 / §11.4.118 / §11.4.123 / §11.4.135. Recommended gate `CM-OPERATOR-ESCAPE-BLUFF-AUDIT` + `CM-COVENANT-114-138-PROPAGATION` (literal `11.4.138`) + paired §1.1 meta-test mutation (gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md §11.4.138. Non-compliance^{11 4 139} is a release blocker. No escape hatch — no `--close-without-bluff-audit`, `--operator-find-is-just-a-bug`, `--skip-permanent-guard` flag.
2026 06 08

§ . . — Fresh-process clean-artifact runtime-signature mandate (User mandate, — —)

Refines §11.4.108. Before any post-deploy validation — ESPECIALLY a non-pixel proxy verification (the subdebug/accessibility-cue channel used for FLAG_SECURE displays) — the harness MUST assert running-artifact == built-artifact: the deploy yielded a CLEAN target (mutable-overlay/userdata wiped) OR a pre-validation check proves no stale overlay shadows the deployed code (e.g. every guarded package — incl. the Presenter that emits the subtitle cue — resolves to the system partition, no per-user override). A stale shadow of the cue-emitting component (e.g. a Presenter APK predating the denylist) makes the proxy report on code that was never deployed — any PASS is a §11.4 PASS-bluff. Each fix declares ONE machine-checkable runtime signature verified on the clean target (the §11.4.108 registry IS the definition of done); for the subtitle class the signature is "the shipped Presenter APK contains the denylist literal (case-insensitive) AND the subdebug channel emits `candidate REJECTED reason=chrome-label` for a menu label."

Classification: universal (§11.4.17). Composes §11.4.46 / §11.4.108 / §11.4.130 / §11.4.135 / §11.4.137. Recommended gate `CM-CLEAN-ARTIFACT-RUNTIME-SIGNATURE` + `CM-COVENANT-114-139-PROPAGATION` (literal 11.4.139) + paired §1.1 meta-test mutation (gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md §11.4.139. Non-compliance^{11 4 140} is a release blocker. No escape hatch — no `--validate-against-running-state`, `--skip-clean-precondition`, `--shadow-OK` flag.
2026 06 09

§ . . — Universal action-prefix system (ACTION_NAME ::) (User mandate, — —)

Forensic anchor — verbatim user mandate (2026-06-09):

"When a user prompt STARTS with `ACTION_NAME ::` (e.g. `BACKGROUND ::` `IMPORTANT: do X...`) the agent MUST replace the `ACTION_NAME ::` prefix with that action's registered expansion text, then execute the rest. For `BACKGROUND ::` the expansion is: 'The following prompt that we will provide MUST BE executed in background in parallel with all main work streams using the subagents-driven development approach! All work done MUST PRODUCE rock solid evidence covered with hard physical proof(s) that all done is working as expected and as specified without any false results and without any bluff!'. The mechanism MUST be UNIVERSAL/extensible — more actions will be added later, each with its own expansion + rules. It MUST work with EVERY CLI agent, be fully decoupled + reusable by every project that includes the constitution submodule, and load + execute out of the box."

Every project under this Constitution MUST support a universal, extensible **action-prefix system**: when a user prompt's FIRST non-blank line starts with an uppercase action token followed by `" :: "` (grammar `^([A-Z][A-Z0-9_]*) ::` — anchored at line start, UPPERCASE-only token, exactly one space on each side of `::`), the agent MUST (1) look the token up in the shared **action registry** (tracked data file `constitution/actions/registry.yaml` , or `$HELIX_ACTION_REGISTRY`); (2) if the token is a registered action, REPLACE the `ACTION_NAME ::` prefix with that action's registered `expansion` text and apply its `rules` ; (3) execute the REMAINDER of the prompt under the expanded instruction. The system is the C-preprocessor expand-then-rescan model applied to prompts (a line-anchored single-token macro): detect → substitute the registered expansion → execute the residual.

Two-layer architecture (both mandatory; LAYER 1 is the universal floor). (LAYER 1 — universal, always-on, out-of-the-box) the action-prefix recognition instruction MUST be mirrored into EVERY agent context carrier the constitution maintains (`CLAUDE.md` , `AGENTS.md` , `QWEN.md` , `GEMINI.md` per §11.4.35) so the agent itself recognises and applies the prefix on every CLI agent — Claude Code, Gemini CLI, Qwen Code, OpenAI Codex CLI, GitHub Copilot CLI, Cursor, Aider, Cline, Continue, Roo Code, and any future agent that reads one of those carriers. (LAYER 2 — mechanical, where the agent exposes a pre-submit seam) a prompt-preprocessing hook reading the SAME registry applies the expansion deterministically (Claude Code `UserPromptSubmit` / `UserPromptExpansion`

hook injecting the expansion via `additionalContext` ; per-agent slash-command equivalents generated from the registry for Gemini/Qwen/Codex). LAYER 2 is the §11.4.109 anti-forgetting upgrade applied to prompt prefixes — the expansion holds even if model recall lapses; LAYER 1 guarantees the system works everywhere with zero extra setup. Honest §11.4.3 boundary (§11.4.6): transparent mechanical free-form `^PREFIX ::` interception is genuinely available ONLY on Claude Code today; on every other agent the free-form form is honoured by LAYER 1 (the agent self-applies) and the slash-command equivalent is the mechanical convenience — this split is documented, never papered over.

Grammar (mandatory, all hold). (1) Anchored at the start of the first non-blank line ONLY; mid-prose tokens never match. (2) UPPERCASE-only token `[A-Z][A-Z0-9_]*` ; lowercase never matches. (3) Separator is exactly `" :: "` (space-colon-colon-space) — avoids C++ `Foo::Bar` , YAML `key: value` , URLs. (4) Stacked prefixes (`A :: B :: rest`) apply outer-to-inner, left-to-right (expand A, rescan, expand B, then the residual is the task). (5) A leading `\` escapes the prefix (literal, no expansion) so action names can be discussed. (6) An unknown token that matches the grammar shape but is NOT registered is NEVER silently expanded or silently dropped — the agent asks which registered action was meant (§11.4.66 + §11.4.105) or treats it literally; it NEVER invents an expansion (§11.4.6). (7) Any prompt not satisfying the grammar is an ordinary prompt; the system is a no-op. The carrier-block GRAMMAR_ADDENDUM (2026-06-09, extended with the arrow form 2026-07-02) additionally recognises FOUR EQUIVALENT forms of every action beyond the bare `::` —

`PREFIX::ACTION_NAME :: <rest>` (namespaced `::`), `/ACTION_NAME <rest>`

(bare slash), `/PREFIX::ACTION_NAME <rest>` (namespaced slash), and

`ACTION_NAME ---> <rest>` (bare arrow, with the namespaced

`PREFIX::ACTION_NAME ---> <rest>` variant) — where `PREFIX` is an action

NAMESPACE (reserved default `DEFAULT`) and the namespace separator `::`

carries NO surrounding spaces (distinct from the action-body `" :: "` and the

arrow-body `" ---> "` , each one ASCII space on either side); the slash, arrow,

and namespaced forms are honoured with the same expansion/execution, the bare

slash form yielding to a colliding built-in/host slash command while the namespaced

slash form is always unambiguous. A leading `\` escapes any of these forms

(`\BACKGROUND :: x` , `\BACKGROUND ---> x` , `\BACKGROUND x`).

Registry is the single source of truth + the extension contract. Adding a new action = adding ONE actions[] row (name , version , summary ≥6 words per §11.4.91, expansion , optional rules + composes_with) — no code change in either layer (the registry is data, not code). Both layers read the same file (§11.4.93-style single-source-of-truth). The first registry entry is BACKGROUND with the operator's verbatim expansion above; it composes §11.4.20/§11.4.70 (subagent-driven), §11.4.58/§11.4.103 (parallel streams), §11.4.89 (background execution), and §11.4.5/§11.4.69/§11.4.107 (captured physical evidence) + §11.4 (anti-bluff). The second registry entry is REMINDER — it re-surfaces previously-scheduled, CRITICAL, status-UNCERTAIN work: the agent FIRST verifies the ACTUAL current status from captured evidence (never assuming done or not-done — §11.4.6 no-guessing), THEN acts on the delta (report the captured proof if genuinely complete, resume from the exact point if partial, action it NOW if not started, or surface the block per §11.4.66 / §11.4.101), ALWAYS producing a status verdict; it composes §11.4.6 / §11.4.87 / §11.4.94 / §11.4.97 / §11.4.103 / §11.4.108 / §11.4.130 / §11.4.147. Three further registry entries are the **severity / handling-priority markers** CRITICAL (highest-priority / potentially release-blocking — maximum urgency + rigor ahead of lower-priority work, track it, auto-activate §11.4.102 systematic-debugging when an issue is involved), IMPORTANT (high-priority, above routine work but below CRITICAL), and NOTE (capture the remainder as durable CONTEXT — request-history §11.4.208 / §11.4.210 + persistent memory when a non-obvious project fact — applied when relevant, not an urgent action unless it contains an explicit action, §11.4.6 record only what the note says); each supports all six §11.4.140 forms incl. the single-colon CRITICAL: / IMPORTANT: / NOTE: , and registering NOTE promotes NOTE: from an ordinary-prose NO-OP to a registered action. They tag HOW the remainder is handled (priority + rigor), not a workable-item Type (§11.4.16). Decoupled + reusable (§11.4.28): the registry + hook + expander carry zero project-specific data and are consumed by reference (the §11.4.80 codegraph_* pattern); a consuming project ships its own registry or inherits the constitution default. Loads out-of-the-box via the §11.4.75 install seam.

Classification: universal (§11.4.17). Composes §11.4 / §11.4.5 / §11.4.6 / §11.4.17 / §11.4.20 / §11.4.28 / §11.4.35 / §11.4.58 / §11.4.66 / §11.4.69 / §11.4.70 / §11.4.75 / §11.4.80 / §11.4.89 / §11.4.91 / §11.4.93 / §11.4.103 / §11.4.105 / §11.4.107 / §11.4.109. Propagation gate `CM-COVENANT-114-140-PROPAGATION` (literal `11.4.140` across the consumer fleet)

- recommended gate `CM-ACTION-PREFIX-SYSTEM` (registry exists + parses + the grammar regex matches/no-matches the canonical fixtures + the LAYER-1 mirror block is present in `CLAUDE.md/AGENTS.md/QWEN.md/GEMINI.md` + the BACKGROUND expansion text matches the registry verbatim + the Claude `UserPromptSubmit` hook + expander script exist and are executable) + paired §1.1 meta-test mutation (corrupt the BACKGROUND expansion in the registry OR strip the `prefix_regex` OR remove the LAYER-1 mirror block from one carrier → the gate FAILs; strip the literal `11.4.140` → the propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.140. Non-compliance is a release blocker. No escape hatch — no `--skip-action-prefix`, `--ignore-prefix`, `--no-registry`, `--invent-expansion-OK`, `--single-layer-only` flag.

§ . . – Token-efficiency mandate (research-derived + operator mandate, – –)

Forensic anchor (operator mandate, 2026-06-09): the AI coding agent's token spend (input AND output) is dominated by an enormous always-loaded governance context — the consumer `CLAUDE.md/AGENTS.md/QWEN.md` plus the constitution submodule (measured ~170,000 tokens of static governance re-sent on EVERY turn: consumer `CLAUDE.md` ~90.5K + `constitution/CLAUDE.md` ~79.6K, a 112-row Applied-Fixes table plus verbatim §11.4.X anchor restatements) — compounded by 350–520K-token subagent transcripts, verbose status, full-file re-reads, and full-file loads where retrieval would do. The same pattern is externally documented as a scaling defect (Claude Code issue #24147 — "Cache read tokens consume 99.93% of usage quota - architectural scaling issue with `CLAUDE.md` re-

reads"). Operator mandate: cut token spend to **30–40% of current (a 60–70% reduction)** WITHOUT degrading quality, performance, or safety, and WITHOUT breaking any existing mechanism.

Every project worked on by AI coding agents **MUST** adopt a token-efficiency regime composed of the following measures, ranked by safety and impact, and **MUST PROVE** the reduction with a measured anti-bluff harness — never an asserted number (§11.4.6 / §11.4.123). The headline 60–70% is the warm-cache target; the rule requires the *measured* best-safe reduction with cited evidence, never the estimate.

The measures (composable, each preserving every existing rule):

- 1. Prompt-cache the static governance prefix (PRIMARY — the single biggest lever).** The always-loaded governance **MUST** form a byte-stable prefix with a cache breakpoint at its end, with nothing volatile rendered ahead of it. Cache reads cost $\sim 0.1\times$ base input price (a 90% discount); writes $1.25\times$ (5-min TTL) / $2\times$ (1-hour). The discount applies to the dominant cost component. **Caching is transparent — the model sees byte-identical input cached or not — so it CANNOT remove a rule, weaken a gate, lower quality, or change a verdict; it changes only billing.** The only failure mode is a *silent invalidator* (timestamp/UUID/unordered-JSON ahead of the breakpoint), which costs money but never corrupts, and is detectable via `cache_read_input_tokens`. The operative discipline: **batch governance edits and keep CLAUDE.md + the constitution byte-stable mid-session** — every governance edit busts the prompt cache (BASELINE.md finding), so all governance changes in a working window are landed together rather than dribbled across turns. Claude Code caches natively; SDK callers use `cache_control:{type:"ephemeral"}`; Gemini/Qwen use their context cache via API-key auth.
- 2. Subagent model-tiering + output-to-file (SECONDARY — biggest non-cache win).** Route MECHANICAL, NON-JUDGMENT subagent work (search, grep, status, doc-export, file-presence, read-only probes) to a cheaper/smaller model (Haiku-class); RESERVE the strong model for ALL reasoning, fix-design (§11.4.102), verdicts, code-review (§11.4.125), and demotion-evidence (§11.4.7). Subagents persist large output to a file and return a pointer instead of a 350–520K-token inline transcript. **The strong model owns every decision and every PASS/FAIL — the cheap model never emits a verdict — so**

§11.4.50 determinism and the anti-bluff covenant are untouched; quality cannot degrade. Bounded by §12.6 (60% memory) and the §11.4.58/§11.4.103 stream caps.

3. **Thin always-loaded INDEX + on-demand detail (progressive disclosure).**

Restructure the consumer governance so the always-loaded file is a concise index (one line per fix / per anchor) and the full bodies are fetched on demand.

Three invariants keep every gate green and every rule reachable: (a) each index line carries the literal `11.4.N` anchor token so the `CM-COVENANT-114-N-PROPAGATION` literal-anchor scans still pass; (b) the canonical full text stays in the tracked, gate-scanned `constitution/Constitution.md` — no rule is deleted, only de-duplicated out of the consumer; (c) the full body is reachable in one hop. This realises §11.4.35's consumer-thin / submodule-canonical split at the byte level — a de-duplication, never a deletion.

4. **Retrieval / CodeGraph-first over full-file loading.** Structural questions (where is X / what calls Y / what would break) → CodeGraph (§11.4.78/§11.4.79, already mandated + installed) or semantic search; region questions → grep-scoped read; never re-read a harness-tracked file. Read-only, already-trusted (§11.4.78) — no behaviour change.

5. **Output-token reduction.** Terse conductor status (no restatement of unchanged governance); `effort:"low"` on the mechanical-subagent allowlist only; structured output where a sink consumes it. Output is ~5× input price, so this is cost-leveraged. Presentation-only — captured-evidence artefacts (§11.4.5/§11.4.69) unchanged.

6. **Tool-call efficiency.** Batch independent tool calls in one block; reuse harness-tracked file state; reuse persisted tool outputs.

7. **Compaction / context-editing for long sessions (defensive).** Prune stale tool-results / old thinking so a 100+ turn cycle doesn't re-pay for an ever-growing transcript; never touches the governance prefix or on-disk evidence.

Mandatory measured proof (anti-bluff — §11.4.6/§11.4.50/§11.4.69/§11.4.123):

the reduction MUST be proven by a token-accounting harness measuring tokens-per-development-cycle BEFORE vs AFTER on a frozen deterministic workload, split into `input_tokens` / `cache_read_input_tokens` / `cache_creation_input_tokens` / `output_tokens` from the authoritative

Anthropic usage object (NEVER tiktoken ; NEVER the client-side total_cost_usd estimate — recompute cost from raw token fields with the published prices), reproduced N times to identical verdict (§11.4.50). Pass = AFTER $\leq 40\%$ of BEFORE (target) OR the measured best-safe reduction with a cited cold-cache reason. The AFTER run MUST ALSO show ZERO regression on the pre-build full sweep, the meta-test mutation sweep (every gate still FAILs its paired mutation — proves M3 did not turn any rule into a bluff), every propagation gate (literal anchors intact), a strong-model reasoning-path probe (proves M2/M5 tiered nothing that decides), and a cache-warm proof (`cache_read_input_tokens > 0`). Cost reduction with quality/safety regression is a §11.4 FAIL, not a win.

No measure may break or degrade any existing mechanism — and the rule is structured so none can: caching is transparent (M1); tiering and low-effort are confined to mechanical non-judgment work (M2/M5); indexing preserves the full rules by reference with literal anchors so the propagation gates pass and the canonical text stays gate-scanned (M3); retrieval uses already-trusted read-only infrastructure (M4). The §11.4 anti-bluff covenant family, §11.4.50 deterministic consistency, §11.4.125 code-review, §11.4.40 full-suite retest, and §12.6 memory ceiling are all unconditionally preserved.

Composes with §11.4.5 / §11.4.6 / §11.4.20 / §11.4.40 / §11.4.50 / §11.4.58 / §11.4.69 / §11.4.70 / §11.4.78 / §11.4.79 / §11.4.80 / §11.4.103 / §11.4.106 / §11.4.123 / §11.4.125 / §11.4.128 / §12.6 / §1.1. Classification: universal (§11.4.17) — the consuming project supplies its model IDs, its mechanical-subagent allowlist, its governance files, its structural-index tool, and its frozen workload per §11.4.35. Propagation gate `CM-COVENANT-114-141-PROPAGATION` (literal 11.4.141 across the consumer fleet) + recommended gate `CM-TOKEN-EFFICIENCY` (the measured BEFORE/AFTER harness exists + produced a verdict $\geq$ target-or-cited-floor + the non-regression sweep is green + cache-warm proof present) + paired §1.1 meta-test mutation (inject a per-turn volatile token ahead of the governance cache breakpoint → `cache_read_input_tokens` collapses → measured reduction falls below the bar → gate FAILs; restore → PASSes — proves the harness measures real caching, not a hardcoded green). Gate-code = separate work item.

Canonical authority: constitution submodule `Constitution.md` §11.4.141. Research + design + measurement: `docs/research/token_efficiency/{RESEARCH, DESIGN, MEASUREMENT, TEST_PLAN}.md` . Non-compliance is a release blocker regardless of context. No escape hatch — no `--skip-token-`

efficiency , --no-cache-governance , --assert-reduction-without-
measuring , --tier-down-reasoning , --inline-all-governance , --
tiktoken-estimate-OK flag exists.

2026 06 09

§ . . — Universal code-review
mandate — every change reviewed, always, no
exception (User mandate, — —)

Forensic anchor — verbatim user mandate (2026-06-09):

"ALL changes we do MUST pass through the code review step!!! ALWAYS!!!"

EVERY change made to ANY repository this Constitution governs — without exception — MUST pass through an INDEPENDENT code-review step BEFORE it is accepted, committed, or built. There is NO change class exempt: source code, fixes, tests, gates, meta-test mutations, documentation, doc-tooling, build/CI scripts, configuration, governance files (Constitution / CLAUDE / AGENTS / QWEN), conductor main-stream edits, sub-agent output, refactors, single-line edits — all of it. This is the ABSOLUTE form of §11.4.125 (which mandated a code-review agent gate after a batch of fixes, before the pre-build sweep + main build): §11.4.142 removes every implicit scoping that §11.4.125's "after all fixes/changes/ implementations are done" phrasing could be read to leave open — there is no "this is just a doc edit" exemption, no "this is just a one-liner" exemption, no "the author already self-reviewed (§11.4.92)" substitution, no "trivial change" carve-out. If a diff exists, it gets an independent review.

Independence is load-bearing. The review MUST be performed by a reviewer structurally separated from the author (a dedicated code-review agent, subagent-driven by default per §11.4.70 / §11.4.20, or a distinct human reviewer) — NOT the author re-reading their own work. §11.4.92's multi-pass self-evaluation is the AUTHOR-side discipline and remains required, but it does NOT satisfy §11.4.142: the whole point of an independent review seam is to catch the blind spots self-review cannot (the structurally-separated reviewer seam §11.4.70 names). A self-review labelled as the code-review step is a §11.4.142 violation of the same class as a §11.4 PASS-bluff at the review layer.

The review is itself anti-bluff. A rubber-stamp "looks good" with no examined findings is not a review (§11.4 / §11.4.1) — the review MUST genuinely analyse the change for correctness, safety (no host §12 / data §9 / target-hardware §11.4.133 regression), will-it-really-work (no solve-A-create-B per §11.4.1), end-user behaviour (§11.4 / §107), and test-genuineness (every covering test exercises the work and catches its negation per §1.1). Its conclusions are captured evidence per §11.4.5 / §11.4.69. **The review iterates to a clean GO per §11.4.134** — any finding (BLOCKING / nit / warning) re-arms the loop; the change is accepted only after a re-run review returns ZERO new findings AND ZERO warnings, with rock-solid physical evidence per §11.4.134.

Composition, not duplication. §11.4.142 is the universal "always, every change, no exception" capstone over the code-review family: §11.4.125 supplies the multi-layer review CONTENT (the six analysis dimensions + the build-gate marker `CM-CODE-REVIEW-GATE-BEFORE-BUILD`); §11.4.134 supplies the iterate-until-GO + rock-solid-evidence LOOP; §11.4.92 supplies the author-side self-review that PRECEDES (never replaces) the independent review; §11.4.70 / §11.4.20 supply the subagent-driven default execution model. §11.4.142 binds them and strips the last exemption: the review step is not a batch-end nicety, it is a precondition on the ACCEPTANCE of any diff. Honest boundary (§11.4.6): a passing review proves the change was independently examined and found sound to the reviewer's evidence — it does NOT replace the §11.4.108 four-layer runtime-signature verification nor the §11.4.40 full-suite retest; it is one of MULTIPLE STRONG LAYERS, and the FIRST one every change crosses.

Forensic FACT (teaching citation, 2026-07-22 — review-before-push is not optional). A transport-hardening change was committed AND pushed to a project's `main` BEFORE any independent review (the author honestly flagged the omission); the post-hoc §11.4.209-substrate review then returned CONDITIONAL/NO-GO with ONE BLOCKING finding — the change had left a security-relevant protection default in the unsafe direction (an operator-keystroke-obscuring option shipped default-off with no disclosure) — plus two Important and four Minor findings, and the remediation consumed a full fix-forward cycle (RED-first test updates per §11.4.115, the launcher fix, three documentation surfaces, and a §11.4.134 re-review) that a pre-push review would have folded into the original change. This is the exact class this anchor forbids: the review found a real security-default defect that was ALREADY public on every mirror when it was found. The

review comes BEFORE commit/push; a post-hoc review that catches a Blocking issue PROVES the mandate, it is not a substitute for it (remediation is fix-forward per §11.4.113 — never a history rewrite).

Classification: universal (§11.4.17) — a platform-neutral acceptance-gate discipline reusable by ANY project; the consuming project supplies its concrete reviewer-dispatch mechanism + the change-acceptance seam (commit wrapper / merge queue / PR gate) per §11.4.35. Composes §11.4.1 / §11.4.4 / §11.4.5 / §11.4.6 / §11.4.20 / §11.4.40 / §11.4.69 / §11.4.70 / §11.4.92 / §11.4.108 / §11.4.110 / §11.4.125 / §11.4.134 / §107 / §1.1. Propagation gate CM-COVENANT-114-142-PROPAGATION (literal 11.4.142 across the consumer fleet) + recommended gate CM-EVERY-CHANGE-REVIEWED (every accepted change carries a fresh independent-review-completed marker for its diff, produced by a reviewer distinct from the author, before acceptance/commit/build) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; accept a diff with no independent-review marker → CM-EVERY-CHANGE-REVIEWED FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md §11.4.142. Non-compliance is a release blocker regardless of context. No escape hatch — no --skip-review, --no-review, --trivial-change-exempt, --doc-edit-exempt, --self-review-suffices, --review-after-commit flag exists.

2026 06 10

§ . . — Real-user-journey mandate
for video-streaming-app full-automation tests
(User mandate, - -)

Forensic anchor — verbatim user mandate (2026-06-10):

"All video streaming apps ... require to choose some title and to press proper UI button to start or resume playing! Proper UI interaction to play exact show with proper content and subtitles is MANDATORY! Without it we just eventually play on 2nd display sample, and that's it mostly! THIS MUST BE ADDED as MANDATORY RULE regarding testing of any video streaming app in general with full automation tests! Universal in root constitution + a consuming project's extensions."

Any full-automation test that asserts a video player / streaming application plays content MUST drive the REAL end-user journey through the app's OWN user interface — launch the app → BROWSE the actual catalog → choose a SPECIFIC title → press the real Play / Resume button → confirm THAT chosen content is genuinely playing, with its correct subtitles, on the intended routing target. A test that bypasses the journey with a sample / demo / built-in-loop clip, a deep-link / `am start -a VIEW` / intent shortcut, a synthetic or pre-staged stream, or any path that does NOT exercise the app's own browse-select-play UI is a §11.4 PASS-bluff at the user-journey layer: it validates ROUTING (that *something* reaches the display) while leaving the user-visible behaviour the operator actually cares about — "the show I picked actually plays" — unproven. Such shortcuts exercise neither the catalog/auth/title-resolution path nor the real Play-button → decoder → routing transition where the end-user defects live (the operator's "we just eventually play on 2nd display sample, and that's it mostly" is exactly this gap).

The mandate (ALL must hold): (1) **Real journey, not a shortcut** — the test drives launch → catalog browse → specific-title selection → real Play/Resume press through the app's own UI (§11.4.48 UI-driven), NEVER a deep-link / intent / sample / loop-clip shortcut; (2) **Chosen-content confirmation** — the PASS proves the SPECIFIC selected title is playing (not merely that pixels move on the target), via the §11.4.107 liveness battery (live advancing frames + frame-advance counter + not-stale-from-previous) on the §11.4.136 real-content path, plus the §11.4.137 subtitle content-correctness oracle for the chosen title's captions (a present cue is NOT a correct cue); (3) **Non-introspectable UIs use the pixel oracle** — when the app's accessibility/semantic hierarchy is blank or unreliable (TV-Compose / leanback / canvas / GL streaming UIs), the test DRIVES input + ASSERTS content via the §11.4.117 CV/OCR pixel oracle (template-match the Play control + ROI-OCR the title/caption), never a hierarchy-only tool; (4) **Login via the credential single source** — apps requiring sign-in authenticate from the project's credential single-source-of-truth (§11.4.10), never hardcoded, never logged; (5) **Honest SKIP, never a faked PASS** — where the autonomous journey is genuinely infeasible (hard human-only login / CAPTCHA, geo-block per §11.4.3, secure-surface pixel-blanking per §11.4.112) the test is `operator_attended SKIP-with-reason` per §11.4.52 + §11.4.3 with a tracked migration item (§11.4.15 / §11.4.16) — NEVER a metadata-only / sample-played / routing-only PASS. Honest boundary (§11.4.6): "the routing fired so the title is playing" is a guess — only the chosen-content

liveness + subtitle oracle on the real journey proves it; a SKIP for genuine infeasibility is correct, a PASS that never pressed the real Play button on the real title is the bluff this anchor forbids.

Classification: universal (§11.4.17) — a platform-neutral end-user-journey discipline reusable by ANY project that ships or integrates video-streaming applications; the consuming project supplies its concrete app roster, login-credential source, routing target, and UI-driving / pixel-oracle harness per §11.4.35. Composes §11.4.48 (UI-driven traversal of the real journey) / §11.4.107 (liveness — the chosen content actually advances) / §11.4.117 (CV/OCR pixel oracle for non-introspectable streaming UIs) / §11.4.136 (real-content end-to-end playback) / §11.4.137 (subtitle content-correctness for the chosen title) / §11.4.52 (autonomous-first; operator-attended only when genuinely infeasible) / §11.4.3 (topology / geo SKIP-with-reason) / §107 (end-user usability is the bar) / §1.1. Propagation gate `CM-COVENANT-114-143-PROPAGATION` (literal `11.4.143` across the consumer fleet) + recommended gate `CM-VIDEO-REAL-JOURNEY-TEST` (every video-streaming-app playback test drives the real browse-select-play journey + confirms the chosen content + subtitles, or SKIPS-with-reason) + paired §1.1 meta-test mutation (replace a real-journey test's browse-select-play path with a deep-link / sample shortcut → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.143. Non-compliance is a release blocker regardless of context. No escape hatch — no `--sample-playback-ok`, `--skip-real-journey`, `--deep-link-suffices`, `--routing-only-pass`, `--no-title-selection` flag exists.

2026 06 10

§ 11.4.143 — Tracked/recorded-device availability-following mandate (User mandate, — —)

Forensic anchor — direct user mandate (2026-06-10): every device the project is tracking / following / recording MUST be availability-followed with the project's already-defined reconnection timings + the sanctioned device-recovery path; a silent recording/tracking gap — papering a disconnect as continuous capture — is a §11.4 bluff.

Every device the project is recording or tracking (test / debug / manual-testing device, across every reachable transport — USB / wireless ADB / SSH / serial / network introspection API) MUST be availability-FOLLOWED: its connection state continuously monitored, and any drop handled — never silently abandoned and never presented as a continuous recording. The always-on recorder under §11.4.128 KNOWS a tracked device is absent (its per-device loop guards on the reachable state) but, lacking a following discipline, merely spins idle — no data captured, no offline event logged, no resume, no escalation — so the recording corpus presents a continuous timeline with a silent hole. That silent hole is a §11.4 PASS-bluff at the recording-integrity layer: it claims continuous capture while no data exists for the gap.

On a tracked device leaving its reachable state the system MUST, automatically: (a) **DETECT** the drop and **log an honest offline event** into the recording corpus (§11.4.6 — a silent gap presented as continuous capture is a fabricated-continuity bluff; §11.4.128 — a silent recording gap defeats always-on recording); (b) **WAIT** for the device to return using the project's ALREADY-DEFINED reconnection timings — never invented numbers (§11.4.6), the SAME grace / reconnect / poll budgets the project's recovery path already uses; (c) **RE-ATTACH** — resume recording / tracking the moment the device returns and log an honest online / resume event; and (d) **ESCALATE** to the project's sanctioned device-recovery path (per §11.4.69 feature class `device_recovery`) if the device does not return within the defined timeout — through the sanctioned, authorization-gated recovery entry point ONLY, never bypassing its gate and never performing a destructive recovery (e.g. a power-cycle) autonomously without that authorization (§11.4.21 / §11.4.101 — high-blast-radius recovery is gated; while blocked the system keeps following the device, logging the blocked-escalation honestly). A tracked-device drop that produces a silent corpus hole, a never-resumed recording, or a never-escalated permanent absence is the bluff this anchor forbids.

Classification: universal (§11.4.17) — a platform-neutral device-following discipline reusable by ANY project that records or tracks devices; the consuming project supplies its concrete tracking transport, the device set, the already-defined reconnection timings, and the sanctioned recovery entry point per §11.4.35. Composes §11.4.128 (always-on recording — §11.4.144 closes its drop-handling gap) / §11.4.69 (`device_recovery` sink-side positive evidence) / §11.4.6 (honest offline / online events, reused-not-invented timings) / §11.4.14 (watchdog children reaped on stop) / §11.4.21 + §11.4.101 (gated, non-autonomous escalation).

Propagation gate `CM-COVENANT-114-144-PROPAGATION` (literal `11.4.144` across the consumer fleet) + recommended gate `CM-DEVICE-AVAILABILITY-FOLLOWED` (the recorder wires a per-device availability watchdog that reuses the defined timings + escalates only through the gated recovery path) + paired §1.1 meta-test mutation (strip the watchdog wiring / let an absence go unlogged → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.144. Non-compliance is a release blocker regardless of context. No escape hatch — no `--skip-availability-following`, `--silent-recording-gap-OK`, `--no-reconnect-wait`, `--invent-reconnect-timing`, `--skip-recovery-escalation`, `--autonomous-power-cycle-OK` flag exists.

2026 06 10

§ . . — Independent multi-angle
impact-research per change (User mandate,
— —)

Forensic anchor — verbatim operator intent (2026-06-10): for EVERY fix / change / new feature, independent agents MUST research potential issues related to that work AND connected/dependent features, from VARIOUS ANGLES, to be SURE no new issues are created, no existing working functionality is broken, and no dangerous code is introduced "like we had with the recents feature", incorporating a VARIETY OF TOOLS.

Forensic anchor (FACT, project-internal): the recents / 3-button-nav path shipped a latent AIDL-interface-contract mismatch that PASSED code review (the diff "looked correct") yet failed at RUNTIME ONLY — it ANR'd on a recents-reaping gesture — because no review angle interrogated the interface contract, the concurrency/lifecycle, or the connected feature it silently broke; the §11.4.108 SOURCE → ARTIFACT → RUNTIME → USER-VISIBLE gap, here caught only after the fact.

§11.4.145 shifts blast-radius discovery LEFT: for EVERY change, INDEPENDENT impact-research agents (subagent-driven §11.4.70/§11.4.20, structurally separate from the author — §11.4.92 self-eval PRECEDES, never satisfies — adversarial "refute-the-change" stance to defeat the documented LLM confirmation-bias failure mode where reviewers accept plausible-but-wrong arguments) research the

change AND its connected/dependent features across a CLOSED SET OF EIGHT ANGLES — (1) **correctness/logic**, (2) **regression** (what existing working feature could break — via call-graph impact enumerating every direct + transitive caller and contract-dependent feature, each mapped to its test), (3) **latent/dangerous-code** — **the recents class** (runtime-only failure, interface/ABI/contract mismatch, concurrency/data-race, resource lifecycle), (4) **security** (taint/injection/secret/unsafe-API), (5) **performance** (latency/memory/thermal under load vs baseline), (6) **host/data/target-hardware safety** per §12/§9/§11.4.133, (7) **cross-feature interaction** (shared state/timing/hardware contention), (8) **business-logic conformance** (matches the spec AND the connected features' contracts, not merely compiles).

Each angle MUST name the TOOL(S) used + obtain the required DEPTH of codebase + business-logic detail (read the diff, the full changed unit, its declared contract, the spec section, and every connected feature — NEVER the diff lines alone) + emit a captured-evidence conclusion per §11.4.5/§11.4.69 (the tool's findings/impact/diff artefact path — "looks fine" is forbidden, §11.4.1). A genuinely-N/A angle is recorded `NOT_APPLICABLE: <reason>` per §11.4.6, never silently skipped. The output is a per-change impact-research REPORT (one section per angle: tool + evidence path + verdict) + a single GO/NO-GO that BLOCKS acceptance/commit/build on ANY unmitigated risk in ANY angle; the change is fixed/mitigated and the affected angles re-researched, iterating to a CLEAN GO (zero unmitigated risk + zero warning, rock-solid physical evidence) per §11.4.134 before the §11.4.125 review gate.

Honest boundary (§11.4.6): a GO proves cross-angle internal-consistency + bounded blast-radius — it does NOT replace §11.4.108 four-layer runtime-signature verification nor the §11.4.40 full-suite retest; it is one of MULTIPLE STRONG LAYERS, the FIRST research pass every change crosses.

Classification: universal (§11.4.17) — the consuming project supplies its concrete toolkit (static-analyzer, call-graph engine, interface-freeze mechanism, behavioral harness) + per-angle depth per §11.4.35. Composes/strengthens §11.4.92 (adds independent tool-backed angle research) / §11.4.125 (its impact-research input) / §11.4.108 (shifts blast-radius discovery left) / §11.4.110 (angles 2/3/7 use the build-graph + interface-freeze) / §11.4.134 (iterate-to-GO) / §11.4.78 / §11.4.79 (call-graph impact engine) / §11.4.107 / §11.4.85 (behavioral angle) / §11.4.133 / §12 / §9 (angle 6) / §11.4.99 (cite latest sources) / §11.4.6 / §1.1. Propagation gate `CM-`

COVENANT-114-145-PROPAGATION (literal 11.4.145 across the consumer fleet) + recommended gate CM-IMPACT-RESEARCH-PER-CHANGE (every accepted change carries a fresh independent impact-research report covering all eight angles with a clean GO, produced by agents distinct from the author, before acceptance/commit/build) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; accept a change with a report missing an angle or with an unmitigated NO-GO → CM-IMPACT-RESEARCH-PER-CHANGE FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md §11.4.145. Non-compliance is a release blocker. No escape hatch — no

```
--skip-impact-research , --single-angle-suffices , --author-researches-own-change , --diff-skim-OK , --no-go-overridable , --trivial-change-no-research flag.
```

2026 06 10

§ . . — Reproduce-first test + same-test-confirms-fix + mandatory extend-to-all-cases workflow (User mandate, — —)

Forensic anchor — verbatim operator intent (2026-06-10): every reported problem MUST be handled by a NAMED three-step test workflow — (1) FIRST author a test that **confirms + reproduces** the reported problem AND is **used as an investigation instrument** to obtain additional forensic data characterising the defect; (2) after the fix, use the **SAME test** to confirm the problem is GONE; (3) THEN **extend** the tests to cover ALL flows, cases, edge cases, and the variety of potential situations for that functionality, confirming no issue of any kind — with every step producing rock-solid PHYSICAL captured evidence, no false results, no bluff.

§11.4.146 does NOT re-author its component disciplines; it NAMES and BINDS them into the operator's exact per-defect sequence and adds ONLY the two emphases the components leave implicit. The mandate (ALL hold):

(STEP 1 — REPRODUCE-FIRST + INVESTIGATE). Before any fix, author the §11.4.43 RED test as a §11.4.115 RED-baseline-on-the-broken-artifact (reproduce the defect on the CURRENT pre-fix artifact, capture defect-present physical evidence per §11.4.5/§11.4.69/§11.4.107). **NEW EMPHASIS (D1):** that same RED

test is ALSO a deliberate **investigation instrument** per §11.4.102(A) — it MUST be used to gather ADDITIONAL forensic data characterising the defect (its triggers, boundaries, input/topology scope, adjacent failure modes), and that characterisation MUST feed BOTH the §11.4.102 root-cause/fix design AND the STEP-3 extend scope. A RED test used only as a binary present/absent gate, with no characterisation captured, satisfies §11.4.115 but NOT §11.4.146's STEP 1.

(STEP 2 — SAME-TEST-CONFIRMS-FIX). After the fix, the SAME test source (the §11.4.115 polarity switch flipped `RED_MODE=1→0`) confirms the defect is ABSENT — RED-on-broken then GREEN-on-fixed, both captured, on a clean target per §11.4.108/§11.4.139, validated-first-after-redeploy per §11.4.130, deterministic per §11.4.50. No separate happy-path test substitutes for the polarity-flipped reproduction test (the §11.4.43/§11.4.115 PASS-bluff).

(STEP 3 — EXTEND-TO-ALL-CASES, mandatory per-fix). NEW EMPHASIS (D2): immediately after STEP 2 confirms the fix, the test MUST be **fanned out across the full case-space of the SAME functionality** — all flows, valid + invalid + boundary (§11.4.85 stress empty/max/off-by-one) + concurrent (§11.4.85 contention) + failure-injection (§11.4.85 chaos) + topology variants (§11.4.3) — confirming no issue of any kind, with PROVABLE enumerated coverage per §11.4.118 (a listed case-set with per-case outcome, never "no other issues found"). This is a REQUIRED step of the per-defect workflow, NOT deferred to a release-cycle discovery sweep and NOT reduced to a single guard. Each fanned-out case that exercises a user-visible path carries rock-solid physical evidence per §11.4.123; each is registered into the §11.4.135 standing regression-guard suite so the functionality's full case-space is permanently guarded; a newly-discovered defect in the fan-out triggers §11.4.4 test-interrupt + re-enters STEP 1 for that new defect.

(ANTI-BLUFF, all steps). Every step's PASS is rock-solid CAPTURED physical evidence per §11.4.123 (§11.4.5/§11.4.69/§11.4.107) — metadata-only / config-only / absence-of-error / grep-without-runtime PASS forbidden (§11.4/§11.4.1); when STEP 1 or STEP 3 validation method is unclear, deep-research-before-declaring-untestable per §11.4.123/§11.4.8/§11.4.99. An operator-found defect the workflow's green suite missed triggers §11.4.138 (bluff-audit + permanent guard).

Honest boundary (§11.4.6): STEP 3 reduces the functionality's unknown-unknown surface; it does not prove zero remaining defects (no finite case-set can, per §11.4.118) — the earned claim is "the reproduce → confirm pair + the enumerated extend case-set are all green with captured proof," un-exercised cases stated as honest gaps, never silently implied clean.

Classification: universal (§11.4.17) — a platform-neutral defect-lifecycle workflow reusable by ANY project; the consuming project supplies its reproduction/polarity harness, case-space enumeration, and capture mechanism per §11.4.35.

Composes §11.4.43 (TDD RED step) / §11.4.115 (RED-baseline + polarity — STEP 1+2 core) / §11.4.102 (systematic-debugging — STEP 1 investigation) / §11.4.130 (validate-fix-first — STEP 2) / §11.4.108 + §11.4.139 (clean-target runtime signature) / §11.4.50 (deterministic) / §11.4.85 (stress+chaos case classes — STEP 3) / §11.4.118 (enumerated discovery coverage — STEP 3) / §11.4.135 (permanent regression guard — STEP 3 registration) / §11.4.3 (topology variants) / §11.4.123 (rock-solid-proof — all steps) / §11.4.5 / §11.4.69 / §11.4.107 (physical evidence) / §11.4.138 (operator-escape) / §11.4.4 (test-interrupt) / §107 / §1.1. Propagation gate `CM-COVENANT-114-146-PROPAGATION` (literal `11.4.146` across the consumer fleet) + recommended gate `CM-REPRODUCE-FIRST-THEN-EXTEND` (every closed defect's fix carries: a §11.4.115 reproduce-first polarity test with captured defect-characterisation [STEP 1], its `RED_MODE=0` GREEN confirmation [STEP 2], AND an enumerated per-functionality extend case-set registered into the §11.4.135 suite [STEP 3] — a closure with the reproduce → confirm pair but NO enumerated extend case-set FAILs) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; close a defect with only the reproduce → confirm pair and no extend-case-set → `CM-REPRODUCE-FIRST-THEN-EXTEND` FAILs; gate-code = separate work item).

(D3) STATUS-CUSTODY mechanical binding — a done-claiming status write is REFUSED without the workflow's machine-written evidence chain (extension, research-derived, 2026-07-17). Forensic FACT (genericised, 2026-07-17): in a consuming project this workflow and §11.4.115 already said exactly this — as prose — and did not bind: of 576 tracked items, 182 claimed a done/ready status and **173 (95%) had NO registered guard at all**; an operator manually sampled 6 done-claiming items and found **6/6 broken** — with 95% unguarded, the *expected draw*, not bad luck; one item recorded 4 Reopened events and 0 terminal events (its status moved with no custody trail), and recurrences re-entered the tracker as NEW ids so even the 52% recorded reopen rate undercounts. A rule that lives only in

prose is re-learned per session and skipped under pressure; a further prose restatement would be the disease's next instance — the deliverable is the BINDING, not more words. Therefore (ALL hold): (1) **the custody chain is a file chain a script checks, with prose nowhere in it** — `item status ∈ done/ready` ⇒ a guard-registry row keyed by the item's EXACT stable id (legacy/alias keys resolved via a checked-in alias map – a guard keyed to a non-item id fails the join and satisfies nothing) ⇒ an executable registered guard ⇒ the §11.4.115(F) machine-written RED+GREEN verdict pair ⇒ evidence files matching the item's feature-class shape ; (2) **Seam A — the status write itself is refused**: the workable-items engine (§11.4.93) refuses a transition to any done/ready status while the item's chain is incomplete; the engine change is project-agnostic per §11.4.28 — it reads a consumer-owned custody configuration, and every project supplies its registry / verdict-store / class-map paths as DATA per §11.4.35; (3) **Seam B — a full-table sweep gate, not a diff check**: on every pre-build run, EVERY done/ready item in the SSoT (minus the ratchet) must hold a complete chain, so an out-of-band raw-DB write that bypasses Seam A is caught at the next sweep rather than never; (4) **Seam C — the release seam** per the §11.4.135 verdict-coverage extension (ABSENCE of a verdict for the release-candidate artifact blocks exactly as a FAIL does); (5) **§11.4.205 construction order** — the FIRST artifact to land is the executable sweep + its golden-bad meta-fixture (a temporary SSoT overlay holding a terminal item with NO custody chain; the meta-test RUNS the sweep against it and requires FAIL — grep-only meta-assertions forbidden per §11.4.201/ §11.4.107(10) — and ALSO asserts the sweep is WIRED: de-wiring its invocation from the pre-build runner flips the meta-test RED, because "installed" is a verified runtime state per §11.4.205(3)); governance text citing the gate lands only in the same commit as the executable, so documenting this clause without implementing the hook is itself caught — the acceptance check is one command: run the sweep; (6) **adoption ratchet** — per the §11.4.135 extension, the pre-existing done-claiming-unguarded backlog is snapshotted once, monotone decrease only. Honest boundary (§11.4.6): the binding proves that no status outruns its evidence and that the REPORTED defect cannot silently return under the guard's conditions — it does NOT prove the feature bug-free (family dedup per §11.4.186 remains), does NOT make a miscalibrated oracle honest (§11.4.107(10) goldens are necessary, not sufficient), does NOT replace the §11.4.185 manual-QA sufficiency gate, and makes target availability a HARD dependency of status progress BY DESIGN — a visible block replaces a silent pass. Recommended gate `CM-STATUS-CUSTODY`

(Seam-A refusal present in the engine + Seam-B full-table sweep wired + ratchet monotone-decreasing) + paired §1.1 mutation (mark a chain-less item done via a raw DB write → the sweep FAILs; de-wire the sweep → its meta-test FAILs; gate-code = separate work item). Composes §11.4.15 / §11.4.28 / §11.4.33 / §11.4.35 / §11.4.54 / §11.4.93 / §11.4.95 / §11.4.106 / §11.4.115(F) / §11.4.135 / §11.4.185 / §11.4.186 / §11.4.201 / §11.4.205.

Canonical authority: constitution submodule Constitution.md §11.4.146. Non-compliance is a release blocker. No escape hatch — no

```
--skip-reproduce-first , --fix-without-red , --skip-extend-to-all-  
cases , --reproduce-confirm-suffices ,  
--defer-extend-to-release-sweep , --single-guard-suffices , --no-  
defect-characterisation flag — and, per (D3), no  
11 4 147  
--status-without-custody , --prose-closure-suffices , --skip-custody-  
sweep , --raw-db-status-write-OK flag.
```

2026 06 10

§ . . — Crashed-agent respawn-until-complete + no-work-loss registry mandate (User mandate, — —)

Forensic anchor — verbatim operator intent (2026-06-10): "any agent that crashed because of something MUST BE respawned and finish its work at some point! We MUST NOT lose any work, forget about or have it corrupted!" Forensic case study (FACT, this session): dispatched subagents died on **transient** causes — API Error: Server is temporarily limiting requests (rate-limit) killed 5 at once, API Error: socket connection closed unexpectedly killed 2, and one left PARTIAL edits (two source files written, the dependent SystemUI sliders + doc + test NOT done). Each was re-dispatched by pure conductor vigilance with **no mechanical guarantee** — the moment conductor context is lost / compacted / the conductor itself crashes, an in-flight-but-dead work unit is silently dropped: work lost / forgotten / corrupted.

Every agent/subagent dispatched to perform a work unit MUST be tracked through its full lifecycle so that a crash NEVER loses, forgets, or corrupts its work. A crashed agent is **NOT** a completed agent (§11.4.6 — "it probably finished before dying" is a forbidden guess; an abnormal termination is positive evidence the work unit is OPEN). The mandate (ALL hold):

(a) Mandatory work-unit / agent REGISTRY. Every dispatched agent/subagent MUST have a durable, append-only, machine-readable registry entry carrying at minimum: a stable agent/work-unit id, its task + file-scope (§11.4.58 PWU manifest), its declared output path(s) + captured-evidence dir (§11.4.5/§11.4.69), its dispatch timestamp, and a status from the CLOSED SET {dispatched | in-flight | crashed | respawned | complete}. The registry reuses the §11.4.116 sync substrate (append-only JSONL event stream + atomically-rewritten status snapshot) — agent-lifecycle transitions are events; the §11.4.116 integrity property holds ("an entry with no dispatch-event cannot show complete "). The registry is the SINGLE SOURCE OF TRUTH for "is this work owed?" — an unregistered dispatch is itself a §11.4.147 violation (an untracked agent cannot be respawned).

(b) Mechanical CRASH-DETECTION + RESPAWN-until-complete. Any abnormal agent termination — transient API/rate-limit error (Server is temporarily limiting requests), socket-close (socket connection closed unexpectedly), any non-completion exit, OR a terminal error after the runtime's own retries are exhausted — MUST flip the registry entry to crashed and keep the work unit OPEN. An OPEN unit MUST be respawned (a fresh agent claims the same id + scope + preserved state) and re-respawned until an agent reaches complete . Respawn is the safe, reversible, bounded-blast-radius decision the conductor takes autonomously per §11.4.101 (never blocks the loop waiting for a human). Transient causes use the project's ALREADY-DEFINED backoff/retry budgets — never invented numbers (§11.4.6); a genuinely-non-transient terminal error (deterministically reproducible after respawn) is investigated per §11.4.102 before further respawn, never spun in a blind retry loop.

(c) PARTIAL-STATE handling — preserve, then resume-or-clean-restart. A crashed agent's uncommitted edits + output doc MUST be PRESERVED (never silently discarded → work lost; never blindly committed → corruption). Before resuming, the respawn MUST run the §11.4.84 quiescence check on the preserved tree (every modified file accounted-for against the unit's scope; no mutation/ // always pass / _mutated_* residue; no half-written/torn artifact). Then EITHER (i) **RESUME idempotently** from the preserved partial when it passes the check (consistent + complete-enough to continue), OR (ii) **CLEAN-RESTART** the unit from a known-good base (§9.2 pre-op backup, reversible per §11.4.101) when the partial is inconsistent / incomplete-unsafely / fails the §11.4.84 check. Either path

guarantees nothing is lost and nothing is corrupted. Per-agent `git worktree` isolation (§11.4.58 L4 / §11.4.84) keeps a crashed unit's partial tree from contaminating other streams.

(d) COMPLETION criterion. A work unit is DONE **only** when an agent reaches `complete` with its required captured evidence/output landed (§11.4.5/§11.4.69 + the §11.4.116 verdict-carries-evidence-path rule). The endless-loop done-condition (§11.4.87 / §11.4.94 / §11.4.97 / §11.4.126) MUST NOT read as satisfied while any registry entry is `dispatched` / `in-flight` / `crashed` / `respawned` -not- yet- `complete` ; the zero-idle survey (§11.4.94) MUST treat every non- `complete` entry as an OPEN actionable item to reclaim. A registry showing `complete` without the landed evidence is a §11.4 PASS-bluff at the agent-lifecycle layer.

Honest boundary (§11.4.6): the registry + respawn guarantee that work is *not lost or corrupted*; it does NOT itself prove the work is *correct* — the respawned unit's output still crosses §11.4.108 four-layer verification, the §11.4.125/§11.4.142 code-review gate, and §11.4.40 full-suite retest. §11.4.147 is the durability layer beneath those, not a substitute. §11.4.147 is the agent-side analogue of §11.4.144 (device availability-following): same detect-drop → wait/backoff → re-attach/respawn → escalate-if-non-transient shape — devices for §11.4.144, dispatched agents for §11.4.147.

Classification: universal (§11.4.17) — the consuming project supplies its concrete registry path/format, its crash-detection signal source (runtime exit codes / sync-stream error events), its respawn-dispatch mechanism, and its already-defined backoff budgets per §11.4.35. Composes §11.4.6 (no-guessing — crash ≠ done) / §11.4.58 (PWU = the work unit + worktree isolation) / §11.4.84 (partial-state quiescence check before resume) / §11.4.87 + §11.4.94 + §11.4.97 + §11.4.126 (loop done-condition + zero-idle reclaim) / §11.4.101 (autonomous respawn-vs-restart decision, reversible-with-backup) / §11.4.116 (registry reuses the sync-channel substrate) / §11.4.102 (non-transient terminal error → systematic-debugging before blind respawn) / §11.4.108 + §11.4.125 + §11.4.142 + §11.4.40 (the respawned output still crosses every downstream verification layer) / §9.2 (clean-restart pre-op backup) / §11.4.128 + §11.4.144 (the recording/device-following siblings — same detect → wait → re-attach → escalate shape, agent-side here). Propagation gate `CM-COVENANT-114-147-PROPAGATION` (literal `11.4.147` across the consumer fleet) + recommended gate `CM-CRASHED-AGENT-RESPAWN-TRACKED` (every dispatched agent has a registry entry; every `crashed` entry is

respawned-until- complete ; the loop done-condition reads no non- complete entries; partial state preserved + §11.4.84-checked before resume) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; mark a crashed entry as the loop's done-condition complete without landed evidence, OR drop a dispatched agent without a registry entry →

CM-CRASHED-AGENT-RESPAWN-TRACKED FAILs; gate-code = separate work item).

(e) API-QUOTA / RATE-LIMIT EXHAUSTION IS A FIRST-CLASS CRASH CLASS (extension, 2026-07-15). Forensic FACT (2026-07-11 + 2026-07-15): background

agents were terminated MID-WORK when the driving subscription hit its weekly / session quota — the agents did not fail, they were KILLED by an EXTERNAL quota boundary, and their in-flight work would have been silently lost had a conductor not manually re-dispatched them. Quota exhaustion MUST be handled by the SAME registry + respawn machinery as any other crash, with these additions: (1) the crash-detection layer MUST record the REAL limit signal + its reason-CLASS per §11.4.196(B) ({session | weekly | subscription} , derived from the captured 429 signature — never faked, never guessed, §11.4.6); (2) the remedy is NOT "wait for the same alias" — the killed unit MUST be REBOUNDED to the next OPERATIONAL alias per §11.4.196(A)/(C) (native-first, provider fallback where configured) and RESPAWNED there; when EVERY alias is cooled the unit bounded-PARKS per §11.4.101 carrying its operational-again epoch, and it stays crashed / OPEN — never complete ; (3) the respawned agent MUST RESUME from the killed agent's captured transcript + on-disk partial state at its EXACT last point (clause (c): preserve → §11.4.84 quiescence check → resume-or-clean-restart), never restart blind, never silently re-do or drop the work; (4) the §11.4.87 / §11.4.94 / §11.4.97 / §11.4.126 loop done-condition MUST NOT read satisfied while any quota-killed entry is non- complete — "the quota ran out" is a PARK, never a DONE. No work may be lost, forgotten, or corrupted by a quota boundary (composes §11.4.196 alias-limit tracking + §11.4.198 always-on multi-alias default — the rebind path is what makes a quota crash recoverable rather than terminal).

Canonical authority: constitution submodule Constitution.md §11.4.147. Non-compliance is a release blocker. No escape hatch — no --skip-agent-registry , --crash-equals-done , --no-respawn , --discard-partial-state , --blind-commit-partial , --forget-dead-agent , --loop-done-with-crashed-entries flag.

§ . 2026 06 10 – Workable-item integrity (status+type+id) + comprehensive structured description + bidirectional external-tracker sync + BLOCKED unblock-choices mandate (User mandate, – –)

Forensic anchor — direct operator mandate (2026-06-10): every workable item MUST carry a VALID status, a VALID type, and a stable unique identifier at all times (never one without all three); every item MUST carry a comprehensive, structured, self-explanatory description (what it is, how it manifests, how to reproduce, acceptance criteria); a BLOCKED item MUST state WHY it is blocked AND the choices that would unblock it; the single source of truth (the SQLite workable-items DB), the rendered docs, and the external tracker (e.g. ClickUp) MUST be regularly, never-missed, fully synchronized in BOTH directions through the mechanical docs-sync engine; the external-tracker push MUST carry statuses, types, assignee, and sub-tasks, default an unset assignee from an env var, and be idempotent.

This anchor BINDS + STRENGTHENS the existing workable-item discipline (§11.4.15 status / §11.4.16 type / §11.4.54 ATM-NNN id / §11.4.21 Operator-blocked / §11.4.91 description-clarity / §11.4.93 + §11.4.95 SQLite SSoT / §11.4.106 docs_chain) into ONE integrity contract spanning the three surfaces — DB ↔ docs ↔ external tracker — and adds the operator's three new emphases (D1 the no-item-without-all-three invariant, D2 the comprehensive structured description, D3 the BLOCKED unblock-CHOICES tightening), citing rather than re-authoring its components. The mandate (ALL hold):

(D1) No item without a valid status + valid type + stable id — across ALL three surfaces. Every workable item, at every moment, MUST carry (i) a `**Status:**` from the §11.4.15 / §11.4.21 / §11.4.90 closed set, (ii) a `**Type:**` from the §11.4.16 closed set {Bug | Feature | Task} , and (iii) a stable, unique, monotonic, append-only [`<ID-NNN>`] identifier per §11.4.54. An item missing ANY of the three in the DB, in the rendered docs, OR in the external tracker is a §11.4.148 violation — the validator FAILs (non-zero exit, release-blocker). The id is the binding key that keeps the same item identifiable across DB rows, doc headings, and tracker tasks; an id present on one surface but absent on another is a sync-integrity failure.

(D2) Comprehensive structured description per item. Every item MUST carry a comprehensive, self-contained, structured description — at minimum WHAT it is (the §11.4.91 ≥6-word / ≥40-char clear meaning), HOW it manifests (observed user-visible / system behaviour), HOW to reproduce (deterministic steps where applicable, §11.4.115 / §11.4.146 reproduce-first), and ACCEPTANCE CRITERIA (the captured-evidence verdict that closes it per §11.4.123 / §11.4.69). The description lives in the §11.4.93 DB `description` column AND renders into the docs AND pushes to the external tracker — never a one-liner stub, never a §11.4.91 anti-pattern fragment. A "fix-it-later"/placeholder description is a §11.4 PASS-bluff at the planning layer (the operator / a future agent / a teammate cannot understand the item).

(D3) BLOCKED items carry WHY + the unblock CHOICES. Tightens §11.4.21: every item whose status is `operator-blocked` (incl. any `Blocked` / `BLOCKED` input alias normalised to the canonical `operator-blocked` value — documented alias, never a silent fork, §11.4.6) MUST carry, beyond the §11.4.21 `operator-Block-Details` WHAT/WHY/WHO, an enumerated set of UNBLOCK CHOICES — the closed list of decisions/actions that would unblock it (e.g. `[A] grant access X` · `[B] approve removal of Y` · `[C] provide credential Z` , mirroring the §11.4.66 2–4-option interactive-clarification shape). The validator FAILs a `BLOCKED` item whose `unblock` field is empty OR carries no enumerated choice marker — "blocked, no way forward stated" is forbidden, the agent / operator must always know the choices that clear the block.

(D4) Bidirectional DB ↔ docs ↔ external-tracker sync, never missed, docs_chain-bound. The §11.4.93 / §11.4.95 git-tracked SQLite DB is the SINGLE SOURCE OF TRUTH; the rendered docs (`Issues` / `Issues_Summary` / `Fixed` / `Fixed_Summary` / per-item docs + their §11.4.65 / §11.4.106 exports) and the external tracker are its DERIVED surfaces. A full synchronization pass MUST run regularly and on every state change so no surface drifts: `DB → docs` (`db-to-md`), `docs → DB` (`md-to-db` , both byte-identical round-trip per §11.4.93), `DB → external-tracker` push, with a drift-proof §11.4.86 fingerprint (sha256 of the sorted item keyset, NOT mtime) gating freshness so a forgotten sync FAILs a pre-build gate. The sync is bound into the §11.4.106 `docs_chain` engine as a registered consumer context (never an ad-hoc parallel script) so drift is mechanically caught, never vigilance-dependent.

(D5) Generic external-tracker sync — statuses/types/assignee/sub-tasks, default-assignee env var, idempotent. The external-tracker push MUST carry: each item's status (collapsed onto the tracker's native status set per §11.4.33 when the tracker's status vocabulary is fixed — the API-cannot-create-statuses case is a §11.4.112 structural constraint, the precise value preserved in a description header, never lost), its type, its assignee (a §11.4.104 participant handle — `created_by` / `assigned_to`), and its sub-tasks. An UNSET assignee defaults from a project-supplied env var (the §11.4.104 operator-username pattern), never hardcoded, never logged (§11.4.10). The push MUST be IDEMPOTENT — a dry-run-then-real, match-by-stable-key (`[<ID>]` name prefix / id custom-field) plan, present⇒UPDATE / absent⇒CREATE, rate-limited, credential-redacted, with a captured sink-side `created=N updated=M failed=0` proof per §11.4.69. The tracker-sync MACHINERY stays project-agnostic (§11.4.28) — the consumer registers its concrete tracker (service, list/board id, field map) at runtime via the public API; the constitution names the discipline, not the vendor.

Anti-bluff (§11.4): every sync run + every validator pass carries captured evidence per §11.4.5 / §11.4.69 (the round-trip diff, the fingerprint match, the tracker `created/updated/failed` line under `qa-results/`); a sync claimed-run with no growing evidence trail, a `complete` / `Fixed` item with no acceptance-criteria description, or a BLOCKED item with no unblock choices are all PASS-bluffs at the workable-item-integrity layer. Honest boundary (§11.4.6): the integrity contract guarantees every item is well-formed + the three surfaces agree, NOT that the item's underlying work is correct — that still crosses §11.4.108 / §11.4.40 / §11.4.123.

Classification: universal (§11.4.17) — the consuming project supplies its concrete DB path, id prefix, external-tracker service + list/board id + field map + default-assignee env var, and docs_chain context per §11.4.35. Composes §11.4.15 / §11.4.16 / §11.4.21 / §11.4.33 / §11.4.34 / §11.4.54 / §11.4.66 / §11.4.86 / §11.4.91 / §11.4.93 / §11.4.95 / §11.4.104 / §11.4.106 / §11.4.112 / §11.4.123 / §11.4.10 / §11.4.28 / §11.4.69 / §11.4.6 / §1.1. Propagation gate `CM-COVENANT-114-148-PROPAGATION` (literal `11.4.148`) + recommended gates `CM-ITEM-INTEGRITY-STATUS-TYPE-ID` (every item carries valid status+type+id on every surface) / `CM-ITEM-COMPREHENSIVE-DESCRIPTION` (every item's description has the WHAT/manifest/repro/acceptance structure) / `CM-BLOCKED-UNBLOCK-CHOICES` (every BLOCKED item enumerates unblock choices) / `CM-TRACKER-SYNC-IDEMPOTENT` (the external-tracker push is bidirectional-fingerprinted + idempotent + docs_chain-

bound) + paired §1.1 meta-test mutations (strip the literal → propagation gate FAILs; drop an item's type / id / description-structure / unblock-choices, OR break the tracker-sync idempotency key → the respective gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.148. Non-compliance is a release blocker. No escape hatch — no `--item-without-status`, `--item-without-type`, `--item-without-id`, `--stub-description-OK`, `--blocked-without-choices`, `--skip-tracker-sync`, `--one-way-sync-OK`, `--non-idempotent-tracker-push`, `--hardcode-assignee` flag.

§ . . — Per-workable-item testing-diary mandate (User mandate, — —)

Forensic anchor — direct operator mandate (2026-06-10): every workable item MUST carry a TESTING DIARY — a chronological, append-only record of every test run against it, each entry capturing date/time, who tested (tested-by), the result, in-depth observations (the technical background a future fix needs), and the action taken + why (whether the run changed the item's status, and the reasoning); both an in-depth diary AND an at-a-glance summary; modelled on the external tracker as a sub-task of the item with a `{TODO | In-progress | Completed}` lifecycle; persisted in the SQLite single source of truth + exported to all required formats + pushed to the external tracker; bound into the mechanical docs-sync engine so it is never missed; minimal-LLM (deterministic tooling, the prose authored by whoever ran the test); under 100% test-coverage discipline.

The testing diary is the evidentiary HISTORY layer for a workable item — distinct from the §11.4.93 `item_history` (which records lifecycle STATE transitions) and from the §11.4.55 `Reopens.md` (which aggregates reopen cycles): the diary records TEST EXECUTIONS, each of which may or may not cause a state transition. The mandate (ALL hold):

(a) Append-only per-item diary table in the SQLite SSoT. A `test_diary` table (additive to the §11.4.93 / §11.4.95 schema, never a competing source) records one row per test run, soft-keyed to the item's stable id, with at minimum: ISO-8601-UTC `date_time` (sortable, deterministic §11.4.50); `tested_by` from a closed set (e.g. `User | Operator | AI-agent | HelixQA`); `result` from the §11.4.45 closed

verdict vocabulary {PASS | FAIL | SKIP} + a short detail; in-depth observations (long markdown — facts or explicit UNCONFIRMED: / PENDING_FORENSICS: per §11.4.6, never "likely"; this is the background a future fix relies on); action_taken + a status_changed flag + from/to (did this run change the item's status, and WHY/why-not); and a §11.4.69 evidence_path + feature_class — with a structural constraint that a PASS row WITHOUT a non-empty evidence path is impossible (a PASS-bluff is rejected by the schema itself).

(b) In-depth diary + at-a-glance summary. Per item, BOTH an in-depth diary document (one section per run, full observations + action + evidence link) AND a derived at-a-glance summary (a rollup VIEW — total/pass/fail/skip runs, last verdict, last run, status-change count, distinct testers, distinct feature-classes — DERIVED, never a duplicate source per §11.4.93). The summary feeds the §11.4.132 risk-ordering inputs (fail-runs / status-changes / last-result) without re-reading every observations blob.

(c) Four-format per-item exports, docs_chain-bound, never missed. Each item's Diary + Diary_Summary render to the project's full export-format set (§11.4.65 + any project-added format) with the freshness invariant (derived-format mtime ≥ source mtime), wired through the §11.4.106 docs_chain engine as a registered consumer context with a §11.4.86 drift-proof fingerprint (sha256 of the sorted diary keyset, NOT mtime) so a diary row added but not exported/pushed FAILs a pre-build gate — the never-missed mechanism.

(d) External-tracker sub-task model with lifecycle. On the external tracker, the workable-item task keeps its description = the item (§11.4.148 D2 — NOT polluted with diary text); each diary run is a CHILD SUB-TASK of that task (type task), carrying a {TODO | In-progress | Completed} lifecycle status (collapsed onto the tracker's native status set per §11.4.33 / §11.4.112 when the tracker's vocabulary is fixed), its observations + action + an Evidence: line (the path, not the raw artefact — §11.4.10 / §11.4.13), and a stable diary-entry idempotency key. The push reuses the §11.4.148 D5 idempotent, rate-limited, credential-redacted, sink-side-proven discipline; the sub-task mapper stays project-agnostic (§11.4.28).

(e) Minimal-LLM, anti-bluff, 100% test-covered. The diary tooling (add / export / tracker-sync / validate) is deterministic bash/Go with ZERO LLM in the data path — the observations prose is authored by whoever ran the test; the tooling only stores/renders/pushes/validates it. The diary is covered per §11.4.27 by every applicable test type (unit + integration-against-the-real-tracker with an honest

§11.4.3 SKIP-when-token-absent, never a faked PASS + export + HelixQA Challenge + paired §1.1 mutation) so a diary PASS-bluff is mechanically impossible. Honest boundary (§11.4.6): the diary guarantees a complete, auditable test-history per item, NOT that any single run's verdict is correct — that rests on the run's own captured evidence per §11.4.69 / §11.4.107 / §11.4.123.

Classification: universal (§11.4.17) — the consuming project supplies its concrete DB path, diary doc layout, export formats, external-tracker sub-task field map, and docs_chain context per §11.4.35. Composes §11.4.6 / §11.4.27 / §11.4.45 / §11.4.50 / §11.4.55 / §11.4.65 / §11.4.69 / §11.4.86 / §11.4.93 / §11.4.95 / §11.4.106 / §11.4.107 / §11.4.123 / §11.4.132 / §11.4.148 / §11.4.10 / §11.4.13 / §11.4.28 / §11.4.33 / §11.4.112 / §1.1. Propagation gate CM-COVENANT-114-149-PROPAGATION (literal 11.4.149) + recommended gates CM-TEST-DIARY-SYNC (the diary fingerprint gate — a diary row not exported/pushed FAILs) / CM-DIARY-PASS-REQUIRES-EVIDENCE (a PASS diary row carries a non-empty evidence path) + paired §1.1 meta-test mutations (strip the literal → propagation gate FAILs; strip the evidence-required constraint OR the sub-task lifecycle / fingerprint → the respective gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md §11.4.149. Non-compliance is a release blocker. No escape hatch — no --skip-testing-diary , --diary-without-evidence , --no-diary-summary , --diary-without-tracker-subtask , --llm-in-diary-data-path , --diary-export-optional flag.

2026 06 11

§ . . — Mandatory deep multi-angle web research per change/issue, before declaring fixed or structural (User mandate, — —)

Forensic anchor — verbatim user mandate (2026-06-11):

"For every single issue we fix or improvement we make — besides tight systematic-debugging, fixing, code review by independent agents, comprehensive tests — we MUST ALWAYS do deep web research from various angles! No matter how big/small/simple/complex, dig deep the internet for articles, technical documentation, APIs, open-source code — EVERYTHING that can help make the best possible solution OR confirm we don't have a more serious problem we're unaware of! We MUST do everything possible to STOP the constant issue-reopening and finally start closing items as fixed+working! Do this ALWAYS in parallel with the main work stream! Add the mandatory rules into the constitution Submodule CONSTITUTION.md, CLAUDE.md, AGENTS.md, QWEN.md and other relevant files. Commit + push to all upstreams. Apply ASAP to every workable item."

Forensic case study (FACT, 2026-06-11). A subtitle-on-secondary-display goal had been verdicted `won't-fix: structurally-impossible` (the §11.4.112 secure-surface pixel-blanking reasoning). A subsequent DEEP multi-angle web research pass OVERTURNED that verdict — it surfaced the real, workable OCR-of-the-PRIMARY-display path (read the caption on the introspectable primary surface rather than the FLAG_SECURE secondary), proving the goal was NOT structurally impossible, only mis-classified for want of research. A single research pass from a different angle converted a "we can't" into a shipped capability. This is the exact reopen-cycle the mandate exists to break: items churn between Fixed ↔ Reopened, or sit mis-labelled `structural / won't-fix`, because the solution-space (or the confirmation that the closure is genuinely safe) was never deep-researched from enough angles.

For EVERY fix / improvement / change / closure — no matter how big, small, simple, or complex — the agent MUST, IN ADDITION to tight systematic-debugging (§11.4.102), the fix itself, independent-agent code review (§11.4.125 / §11.4.142), and comprehensive multi-layer tests (§11.4.4(b) / §11.4.40), perform a DOCUMENTED **deep multi-angle web research pass** that digs the internet from VARIOUS angles — articles, official + vendor technical documentation, API references, standards, issue trackers, maintainer guidance, reusable open-source code — to BOTH (i) discover the best possible solution AND (ii) confirm there is

NOT a more serious underlying problem the team is unaware of. This BINDS + STRENGTHENS the existing research/proof anchors into one mandatory, reopen-breaking, run-in-parallel discipline; it does not re-author them:

- **(A) No closure-as-fixed/structural WITHOUT a documented deep-research pass.** No issue may be marked `Fixed / Implemented / Completed` (§11.4.33) OR classified `structurally-impossible won't-fix` (§11.4.112) UNTIL a deep multi-angle research pass is documented for it. §11.4.8 already requires deep research before designing a non-trivial fix and §11.4.123 already requires it when validation is unclear; §11.4.150 makes the pass UNCONDITIONAL — required for EVERY workable item, however trivial, at the closure AND structural-verdict gates — precisely because the forensic case proved a "trivial-looking" structural verdict was wrong.
- **(B) Multiple angles, not a single lookup.** The pass MUST interrogate the problem from several distinct angles (e.g. official-docs / standards / known-bug-trackers / alternative-approach / failure-mode / security / performance / platform-constraint / open-source-precedent — the consuming project picks the applicable subset, ≥ 2 genuinely-distinct angles), so a single-source confirmation-bias miss (the §11.4.145 LLM-confirmation-bias failure mode) cannot pass as research. A one-link drive-by is NOT a deep multi-angle pass.
- **(C) Confirm-no-bigger-problem, not just find-a-fix.** The pass MUST explicitly seek evidence that the chosen fix does NOT mask a deeper defect AND that the closure (or the `structural` verdict) is genuinely safe — the "do we actually have a more serious problem we're unaware of?" check. Absence of a found bigger problem is recorded as a positive finding only when the search that would have surfaced it is enumerated (composes §11.4.118 discovery-completeness — "we found nothing worse" requires "here is what we searched").
- **(D) Latest-source + cited.** Sources MUST be the LATEST authoritative versions per §11.4.99 (never training-data / memory), each cited by URL + access date in the item's research artefact AND the closure commit footer (`Deep-research <date>: <urls>` OR the literal `NO external solution found – original work` per §11.4.8 when the search genuinely yields nothing). Stale or memory-sourced "research" is a §11.4.99 + §11.4.150 violation.

- **(E) Reopen-breaking is the PURPOSE.** The explicit goal is to STOP the constant Fixed ↔ Reopened churn and finally close items as fixed-AND-working: a reopen (§11.4.34 / §11.4.55) whose root cause a deep multi-angle pass would have surfaced is a §11.4.150 miss; and a `structural / won't-fix` verdict reached without the pass is forbidden (the case study's overturned verdict is the canonical anti-pattern). Composes §11.4.7 (demotion-evidence — a reopen IS a demotion, and the research is part of the demotion-prevention) + §11.4.112 (structural-impossibility now REQUIRES the deep-research pass + its cited authorities before the verdict is earned).
- **(F) ALWAYS in parallel with the main stream.** The research pass MUST run as a background, subagent-driven work stream (per §11.4.70 / §11.4.20 / §11.4.103 continuous-parallel-stream + §11.4.89 background execution) concurrent with the main fix/build/test work — it NEVER serialises the main stream or stalls the loop (§11.4.94 / §11.4.97 / §11.4.126). A blocked research angle parks that angle, it does not pause the closure pipeline (§11.4.101).
- **(G) Apply ASAP to every workable item.** The discipline applies retroactively + going forward to EVERY workable item (§11.4.93 / §11.4.95 SSoT); items closed without a documented pass are re-audited as part of the §11.4.40 / §11.4.42 release-gate sweep.

Honest boundary (§11.4.6): a deep-research pass reduces the unknown-unknown surface and breaks the most common reopen causes — it does NOT prove zero remaining defects (§11.4.118) and does NOT replace §11.4.108 four-layer runtime-signature verification, §11.4.125 / §11.4.142 independent review, or §11.4.40 full-suite retest. It is one of MULTIPLE STRONG LAYERS — the research layer that every fix/closure/structural-verdict additionally crosses. "We probably don't have a bigger problem" without the enumerated multi-angle search is a guess (§11.4.6), never a finding.

Classification: universal (§11.4.17) — a platform-neutral research-and-closure discipline reusable by ANY project; the consuming project supplies its concrete research corpora, angle set, item tracker, and closure-commit-footer convention per §11.4.35. Composes §11.4.8 (deep-web-research-before-implementation — §11.4.150 makes it unconditional + multi-angle + closure-gating) / §11.4.99 (latest-source verification of the cited authorities) / §11.4.123 (rock-solid-proof-or-deep-research — §11.4.150 raises "when unsure" to "always") / §11.4.118 (discovery-completeness — the confirm-no-bigger-problem search is enumerated) / §11.4.145

(independent multi-angle impact-research — §11.4.150 is its external-research sibling, adversarial against confirmation bias) / §11.4.125 / §11.4.142 (independent review — research precedes, never substitutes) / §11.4.7 (demotion/reopen evidence) / §11.4.112 (structural-impossibility now requires the cited-authorities pass) / §11.4.34 / §11.4.55 (reopen attribution — the churn this breaks) / §11.4.70 / §11.4.20 / §11.4.89 / §11.4.103 (parallel background subagent stream) / §11.4.40 / §11.4.42 (release-gate re-audit) / §11.4.93 / §11.4.95 (per-item SSoT). Propagation gate `CM-COVENANT-114-150-PROPAGATION` (literal `11.4.150`) + recommended gate `CM-DEEP-RESEARCH-PER-ISSUE` (every closed/structural-verdicted item carries a documented multi-angle deep-research artefact + cited-source closure footer, run in parallel, before the closure/structural verdict is accepted) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; close an item or reach a `structural` verdict with no documented multi-angle research artefact / cited footer → `CM-DEEP-RESEARCH-PER-ISSUE` FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.150. Non-compliance is a release blocker. No escape hatch — no `--skip-deep-research`, `--single-source-suffices`, `--trivial-change-no-research`, `--close-without-research`, `--structural-without-research`, `--serialise-research`, `--research-from-memory-OK` flag.

§ . . — Project-prefixed release-tag/version-naming mandate (User mandate, - -)

Forensic anchor — verbatim user mandate (2026-06-12):

"Every release tag and version name we create — on the main repository and on every Submodule we own — MUST be prefixed with the project's release prefix, e.g. `myproject-1.0.0-dev-0.0.1`. The prefix MUST come from `HELIX_RELEASE_PREFIX` in our `.env` if it is set, otherwise from the lowercased project root directory name. The SAME prefix MUST be used across the main repo and all owned Submodules in one release so a release is greppable across every repository."

Every release tag AND every version name created on the main repository AND on every owned-by-us submodule (per §11.4.28) MUST be prefixed with the project's release prefix, in the form `<PREFIX>-<version>` (canonical example:

`myproject-1.0.0-dev-0.0.1`). The prefix makes a release identifiable + greppable across every repository the release spans, so a single `git tag -l '<PREFIX>-*'` (or equivalent) enumerates the whole release surface in one query.

Prefix resolution order (closed-set, deterministic — §11.4.6 no-guessing):

1. **HELIX_RELEASE_PREFIX** from the project's `.env` — authoritative when set. `.env` is git-ignored per §11.4.30; the variable MUST be documented in the tracked `.env.example` placeholder so a fresh clone knows the contract (the regeneration-mechanism discipline of §11.4.77 applies — `.env` is re-obtained, never committed).
2. **Fallback = the lowercased snake_case form of the project root directory name** (no spaces), per §11.4.29 — used whenever `HELIX_RELEASE_PREFIX` is unset/empty. The fallback is deterministic from the checkout, so a release prefix is always resolvable without operator input.

The prefix MUST be IDENTICAL across the main repo and all owned

submodules within a single release — a release that tags the main repo

`<PREFIX>-1.2.0` while tagging an owned submodule `<OTHER>-1.2.0` (or unprefixed) is a §11.4.151 violation: the cross-repo grep no longer enumerates the release. Version codes increment monotonically within the prefix (`<PREFIX>-...-0.0.1` → `<PREFIX>-...-0.0.2` → ...), never reset, never skipped — same append-only discipline as §11.4.54's ticket ids applied to the version axis.

Honest boundary (§11.4.6): the prefix guarantees a release is identifiable + uniform across every repository, NOT that the release contents are correct — the tag is still only created after the §11.4.40 full-suite retest GREEN and reaches every upstream via the §11.4.113 merge-onto-latest-main path (NEVER a force-push), fanned out per §2.1.

Classification: universal (§11.4.17) — a platform-neutral release-naming discipline reusable by ANY project; the consuming project supplies its concrete prefix value + the `HELIX_RELEASE_PREFIX` env var per §11.4.35. Composes §2.1 (multi-upstream push — the prefixed tag fans out to every remote) / §11.4.29 (lowercase snake_case — the fallback prefix follows it) / §11.4.30 (`.env` git-ignored — the authoritative prefix source) / §11.4.40 (full-suite retest before the tag is created) /

§11.4.113 (absolute no-force-push — the prefixed tag reaches upstreams via merge-onto-latest-main) / §11.4.126 (the release-scope terminal condition is a published, prefixed tag). Propagation gate `CM-COVENANT-114-151-PROPAGATION` (literal `11.4.151` across the consumer fleet) + recommended gate `CM-RELEASE-PREFIX-NAMING` (every release tag/version on the main repo + every owned submodule carries the resolved `<PREFIX>-` prefix, identical across the release) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; create an unprefixed or differing-prefix release tag → `CM-RELEASE-PREFIX-NAMING` FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.151. Non-compliance is a release blocker. No escape hatch — no `--no-release-prefix`, `--unprefixed-tag`, `--prefix-optional`, `--differing-submodule-prefix` flag.

2026 06 13

§ . . — Crashlytics-recorded-data continuous monitoring + systematic-debug + regression-test-coverage mandate (User mandate, - -)

Forensic anchor — verbatim user mandate (2026-06-13):

"For every project that has Firebase Crashlytics enabled / wired, we MUST continuously monitor ALL of the Crashlytics-recorded data — crashes, ANRs, performance traces, and non-fatals — systematically debug each, fix and improve, and cover everything with validation and verification tests. This MUST be checked regularly, with no false results and no bluff of any kind!"

Every project that has Firebase Crashlytics enabled / wired (the SDK linked into a shipping artifact, crash + non-fatal + ANR reporting active) MUST treat the Crashlytics console as a first-class captured-evidence channel from real end-user devices and continuously process every datum it records. Crashlytics IS the user-visible-behaviour signal §11.4 exists to protect: a recorded crash / ANR / performance regression / non-fatal is a real end user hitting a broken feature, and a recorded issue left unmonitored, undebugged, unfixed, or uncovered-by-test is the

precise "tests pass while the feature is broken in the wild" failure mode the §11.4 covenant forbids — an open Crashlytics issue with a green test suite is a §11.4 PASS-bluff at the field-telemetry layer.

§11.4.152 STRENGTHENS + COMPLETES §11.4.47 (Firebase data review mandate): §11.4.47 owns the periodic REVIEW + dedup + Issue-creation pass that surfaces Crashlytics/Analytics/Performance findings and maps them to tracker entries; §11.4.152 owns what happens to each surfaced item AFTER it is recorded — the systematic-debug → fix/improve → regression-test-coverage lifecycle that drives it to a proven, regression-immune closure. §11.4.47 finds it; §11.4.152 fixes it and proves it stays fixed. Both are mandatory; neither substitutes for the other.

The mandate (ALL must hold):

1. **Continuous monitoring of ALL four Crashlytics surfaces.** The monitoring pass MUST cover fatal crashes, ANRs, performance traces / regressions, AND non-fatals — all four. Skipping any one surface is a §11.4 PASS-bluff (a regression surfaces on the skipped axis and the team never sees it — the §11.4.47 three-source lesson extended to the non-fatal + ANR + performance axes). Non-fatals in particular are the silent class: every catch / fallback / recovered-error path that records a non-fatal is a real degraded user experience and MUST be triaged, not ignored because the app did not crash.
2. **Systematic-debugging of each recorded issue (reproduce-before-fix).** Every Crashlytics-recorded issue selected for fix MUST go through the §11.4.102 systematic-debugging arc — Iron Law: NO FIX WITHOUT ROOT-CAUSE INVESTIGATION FIRST — root-cause → pattern → falsifiable hypothesis proven with captured evidence → fix designed against the proven cause. The §11.4.115 RED-baseline-on-the-broken-artifact reproduces the recorded defect (the recorded stacktrace / ANR trace / non-fatal context IS the §11.4.5/§11.4.69 captured evidence of defect-present) BEFORE the fix; a fix without a prior reproducing test is a §11.4.43/§11.4.123 violation (no falsifiability evidence the fix addresses the root cause vs masks the symptom). Where the recorded issue cannot be reproduced, the deep-research-before-declaring-untestable path (§11.4.123 / §11.4.150) is mandatory before any "cannot reproduce" classification.

3. **A fix / improvement for every confirmed issue.** Every confirmed Crashlytics issue (severity per the §11.4.47 classification table) MUST receive a real fix or improvement landed against its proven root cause (§11.4.9 / §11.4.43 / §11.4.108 four-layer verification on a clean artifact). "Acknowledged in the console" is not a fix; muting / ignoring a recurring issue without a tracked rationale is a §11.4.6 / §11.4.90 violation.
4. **Validation + verification regression-test coverage per closed issue.** Every Crashlytics issue closed by a fix MUST, in the SAME commit as the fix (§11.4.43 DOCUMENT step), register a permanent §11.4.135 regression guard into the standing suite — a §11.4.115 polarity test whose `RED_MODE=1` captures the recorded defect on the pre-fix artifact and whose `RED_MODE=0` is the standing GREEN guard asserting the defect is ABSENT. The guard MUST exercise the same user-reachable code path the Crashlytics stacktrace / trace identifies, with rock-solid captured evidence per §11.4.5 / §11.4.69 / §11.4.107 / §11.4.123. A Crashlytics issue marked resolved WITHOUT a falsifiable regression test is FORBIDDEN — it is the canonical recurrence vector (the issue silently returns months later, invisible to the green suite), the §11.4.138 operator-escape class applied to field telemetry. Each closure carries the §11.4.34 demotion-evidence + a closure-log entry recording the Crashlytics issue id / console URL, the root-cause analysis, the fix commit, and the validation + verification test paths.
5. **Regular cadence — checked repeatedly, never once.** The monitoring + debug + fix + cover cycle MUST fire on a regular cadence, reusing the §11.4.47 five-trigger set (pre-build / pre-flash / pre-distribute / pre-tag blocking; daily + post-deployment-burn-in non-blocking) so newly-recorded issues are caught while they are cheap. A one-time sweep that is never repeated is a §11.4.152 violation; field telemetry accrues continuously and so must its processing.
6. **No false results, no bluff.** Every verdict in the cycle is captured evidence per §11.4 / §11.4.1 / §11.4.6 — a "no new Crashlytics issues" claim requires the enumerated monitored-surfaces + queried-window as proof (§11.4.118 absence-of-evidence is not evidence-of-absence); a "fixed" claim requires the RED → GREEN flip on the artifact (§11.4.115 / §11.4.130); a regression guard whose paired §1.1 mutation does not make it FAIL is itself a bluff gate.

Honest boundary (§11.4.6): processing every recorded issue reduces the known-field-defect surface and breaks the silent-recurrence vector — it does NOT prove zero remaining field defects (a defect no user has hit yet is not yet recorded), and it

does NOT replace §11.4.108 runtime-signature verification or §11.4.40 full-suite retest. It is the field-telemetry-driven layer that complements them; an issue is "closed" only when its regression guard is GREEN on a clean artifact, never when the console mark is flipped.

Classification: universal (§11.4.17) — a platform-neutral field-telemetry-processing discipline reusable by ANY project that wires Firebase Crashlytics (or, by the same shape, any equivalent crash/ANR/perf/non-fatal field-reporting backend); the consuming project supplies its concrete Crashlytics project handle, console-access credential (§11.4.10, never logged), severity table, and per-issue closure-log path per §11.4.35. The reference project-level instantiation is a consuming project's §6.O (Crashlytics-Resolved Issue Coverage Mandate — per-issue validation test + Challenge test +

`.lava-ci-evidence/crashlytics-resolved/<date>-<slug>.md` closure log) + §6.AC (Comprehensive Non-Fatal Telemetry Mandate — every catch / fallback / recovered-error path records a non-fatal with triage context), which together are the project-specific embodiment of this universal clause. Composes §11.4 / §11.4.1 / §11.4.5 / §11.4.6 / §11.4.9 / §11.4.34 / §11.4.40 / §11.4.43 / §11.4.47 / §11.4.69 / §11.4.90 / §11.4.102 / §11.4.107 / §11.4.108 / §11.4.115 / §11.4.118 / §11.4.123 / §11.4.130 / §11.4.135 / §11.4.138 / §11.4.150 / §1.1. Propagation gate `CM-COVENANT-114-152-PROPAGATION` (literal `11.4.152` across the consumer fleet) + recommended gate `CM-CRASHLYTICS-ISSUE-FULLY-COVERED` (every closed Crashlytics issue carries a systematic-debug root-cause record + a registered §11.4.135 regression guard + a closure-log entry citing the console issue id/URL + the validation/verification test paths; a console-resolved issue with no falsifiable regression guard FAILs) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; mark a Crashlytics issue resolved with no registered regression guard / no closure log → `CM-CRASHLYTICS-ISSUE-FULLY-COVERED` FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.152. Non-compliance is a release blocker. No escape hatch — no `--skip-crashlytics-monitoring`, `--monitor-crashes-only`, `--skip-non-fatals`, `--resolve-without-regression-test`, `--console-mark-is-fixed`, `--monitor-once`, `--mute-without-rationale` flag.

§ . . — Comprehensive per-feature Status + Status_Summary document set with mandatory video-recording confirmation (User mandate, — —)

Forensic anchor — verbatim user mandate (2026-06-15): "We MUST CREATE proper Status and Status_Summary documents under docs/features containing and covering all system components, all client apps, and all features we have and we have ported from all cli_agents! ... in-depth with no a single feature left ... table ... always up to date and in sync ... status of its usability implementation tests coverage and the confirmation of practical use confirmed on video recording with no false results or bluff of any kind! ... tightly connected using docs_chain Submodule ... Existence of such Status (with Status_Summary) document with all exported files (PDF, DOCX, HTML) is mandatory for every single project."

Every project MUST maintain, under docs/features/ , a comprehensive **feature Status document set** (Status.md + its §11.4.56 Status_Summary.md companion) that enumerates EVERY system component, EVERY client application/binary/surface (TUI / CLI / Web / desktop / mobile / API / gRPC / library / submodule / infrastructure), and EVERY feature — including features ported in from any incorporated CLI-agent / submodule catalogue (§11.4.74) — with NO single feature left out. The mandate (ALL hold):

(1) **Total, categorized feature coverage** — a per-feature table covering the FULL feature surface, organized Component → Category, reconciled against the actual codebase (a feature present in code but absent from the table, or a table row with no code, is a §11.4.153 violation; "looks complete" without a code-reconciliation pass is a §11.4.6/§11.4.118 bluff).

(2) **Per-feature fields** — each feature row MUST carry: Component, Feature, Category, Implementation status, Wiring status (genuinely reachable by an end user, not merely compiled — §11.4.108), Real-use status, Tests-coverage status (the four-layer §11.4.4(b) coverage state), Validation status (closed vocabulary PASS / FAIL / SKIP / PENDING_FORENSICS / OPERATOR-BLOCKED per §11.4.45), and **Video-recording confirmation** (a path to the captured real-use video proving the feature works for the end user, or an honest gap marker).

(3) **Mandatory per-feature real-use video confirmation** — each user-visible feature's "confirmed working" claim MUST be backed by a recorded real-use video showing a genuine end-user scenario: real prompts → real LLM/service responses producing real results, NEVER a frozen/stale frame (§11.4.107 liveness), NEVER a faked, mocked, demo-loop, or bluff response, NEVER an LLM error/garbage response passed off as success. Videos are captured evidence per §11.4.2/§11.4.5 and stored at the project-declared recording path (§11.4.35); a "confirmed" row with no real video, or a video whose content is a bluff, is a §11.4 PASS-bluff at the feature-status layer. Where autonomous video capture is genuinely infeasible the row is an honest §11.4.3/§11.4.52 SKIP-with-reason + tracked migration item — NEVER a faked confirmation.

(4) **Mandatory video-analysis remediation loop** — every recorded video MUST be analysed (its presented data recognized + checked); any defect the video surfaces MUST trigger the §11.4.102 systematic-debugging → fix → §11.4.146 retest → re-record loop, iterating to a clean GO per §11.4.134, before the feature is marked confirmed. A video that exposes a broken feature is a §11.4.4 test-interrupt, not a confirmation.

(5) **Always-in-sync, mechanically enforced** — the feature Status set is a §11.4.45-class roster/corpus-backed Status doc: kept in sync via the §11.4.106 docs_chain engine + a §11.4.86 drift-proof fingerprint (sha256 of the sorted feature-key roster AND the sorted video-artefact roster, NOT mtime) so that adding/removing/renaming a feature OR a video artefact mechanically re-syncs the docs + exports out-of-the-box; stale = violation.

(6) **Four-format export** — the feature Status set MUST export to HTML + PDF + DOCX (this doc class ADDS DOCX to the §11.4.65 universal HTML+PDF export set; the other doc classes are unchanged), all kept in sync per §11.4.60.

(7) **Follows the Status-doc family** — composes §11.4.44 (revision header), §11.4.45 (Status maintenance + captured-evidence table + operator-blocked-at-top), §11.4.56 (Status_Summary two-audience parity), §11.4.57 (README doc-link row), §11.4.59/§11.4.60 (always-sync composite).

(8) **MP4 format REQUIRED.** All video confirmations MUST be in `.mp4` format (H.264, `movflags +faststart` , `pix_fmt yuv420p`). Window-specific capture ONLY (§11.4.159(A)). Vision validation REQUIRED (§11.4.159(D)). `.cast` (asciinema) files are supplementary only — the primary evidence is the `.mp4` video.

Honest boundary (§11.4.6): the feature Status set guarantees a complete, video-confirmed, always-synced ledger of what works for the end user — it does NOT itself prove correctness of any single feature (that rests on the feature's own §11.4.69/§11.4.107/§11.4.123 captured evidence the row cites), and it does NOT replace §11.4.40 full-suite retest. Classification: universal (§11.4.17) — the consuming project supplies its concrete `docs/features/` layout, feature roster source, recording path, and DOCX exporter per §11.4.35. Composes §11.4.2 / §11.4.5 / §11.4.44 / §11.4.45 / §11.4.52 / §11.4.56 / §11.4.57 / §11.4.59 / §11.4.60 / §11.4.65 / §11.4.86 / §11.4.102 / §11.4.106 / §11.4.107 / §11.4.108 / §11.4.118 / §11.4.123 / §11.4.134 / §11.4.146 / §1.1. Propagation gate `CM-COVENANT-114-153-PROPAGATION` (literal `11.4.153`) + recommended gates `CM-FEATURE-STATUS-COMPLETE` (every code-present feature has a Status row + every row maps to code) + `CM-FEATURE-STATUS-VIDEO-CONFIRMED` (every user-visible confirmed row cites a real captured video or an honest SKIP) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; mark a feature confirmed with no real video / drop a code-present feature from the table → the feature-status gates FAIL; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.153. Non-compliance is a release blocker. No escape hatch — no `--skip-feature-status` , `--feature-without-video` , `--frozen-video-OK` , `--bluff-response-video-OK` , `--skip-video-analysis` , `--feature-ledger-incomplete-OK` , `--no-docx-export` , `--allow-feature-ledger-drift` flag.

§ . . — Window-scoped capture + fresh-corpus rotation for feature/QA recordings (User mandate, — —)

Forensic anchor — verbatim user mandate (2026-06-15): "IMPORTANT: Make sure that recording you perform does record the window containing apps and services, not the whole desktop or monitor screen!" + "NOTE: All old recording files MUST BE removed when new one starts!"

Refines/strengthens the §11.4.2/§11.4.5/§11.4.107/§11.4.153 recording disciplines with two capture-hygiene invariants every feature/QA video MUST satisfy. **(A)**

Window-scoped capture, NOT whole-screen. A recording that proves a feature works MUST capture ONLY the window / surface / viewport of the application or service under test — the app's own window (GUI), its terminal pane (CLI/TUI), its browser tab/viewport (web), or its device/emulator/simulator frame (mobile) — NEVER the whole desktop, the whole monitor, or unrelated windows. Whole-desktop capture is forbidden because it (i) leaks unrelated/operator-private content into a committed artefact (a §11.4.10/§11.4.83 hygiene breach), (ii) dilutes the §11.4.107 liveness/freeze/frame-advance oracle with non-feature pixels (a defect can hide in the noise — a §11.4 evidence-integrity weakening), and (iii) makes the §11.4.137-class OCR/ROI content oracle unreliable. The capture mechanism MUST target the window/region by stable identity (window id/title, device serial, browser context, tmux/terminal target) per §11.4.111 — never a fixed full-screen device index that happens to contain the window. Where the platform genuinely cannot capture below whole-screen for the surface under test, that is an honest §11.4.3 SKIP-with-reason + tracked migration item, NEVER a whole-screen capture passed off as window-scoped. **(B) Fresh-corpus rotation — remove old recordings when a new run starts.** When a new recording run for a given scope begins, the prior run's stale recording files for that scope MUST be removed FIRST, so the live recording corpus always reflects the CURRENT artefact/run and a reader never mistakes a stale video for current evidence (the §11.4.107 not-stale-from-previous principle applied to the corpus, and the §11.4.86 roster-freshness principle applied to video artefacts). Honest boundary (§11.4.6 + §9.2): "remove old" means the agent's OWN prior recordings for the SAME scope/project at the project-declared recording path — NEVER another project's, another scope's, or operator-authored files the agent did not create; when uncertain whether a file is in-scope-stale vs foreign, do NOT delete it — surface it (§11.4.122). Curated evidence already committed under docs/qa/<run-id>/ (§11.4.83) is the durable record

and is NOT deleted by rotation; rotation operates on the raw working recording path (§11.4.128 raw corpus), not the committed evidence trail. Classification: universal (§11.4.17) — the consuming project supplies its concrete window-capture mechanism (per surface class), recording path, and scope-key per §11.4.35. Composes §11.4.2 / §11.4.5 / §11.4.10 / §11.4.83 / §11.4.86 / §11.4.107 / §11.4.111 / §11.4.122 / §11.4.128 / §11.4.137 / §11.4.153 / §9.2 / §11.4.6. Propagation gate `CM-COVENANT-114-154-PROPAGATION` (literal `11.4.154`) + recommended gate `CM-WINDOW-SCOPED-FRESH-CORPUS-RECORDING` (every feature/QA recording is window/surface-scoped not whole-screen + a new run removes its own prior in-scope recordings first) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; a whole-screen capture where window-scoped is feasible, or a new run that leaves its own stale prior recordings, → the recording gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.154. Non-compliance is a release blocker. No escape hatch — no `--whole-screen-capture-OK`, `--skip-window-scope`, `--keep-stale-recordings`, `--no-corpus-rotation`, `--full-desktop-recording` flag.

(C) MP4 auto-conversion REQUIRED. Any `.cast` file produced MUST be auto-converted to `.mp4` via `agg + ffmpeg` immediately after capture. The `.mp4` is the primary evidence; `.cast` is supplementary only. Conversion command: `agg input.cast output.mp4 --renderer-res 1920x1080` followed by `ffmpeg -i output.mp4 -c:v libx264 -pix_fmt yuv420p -movflags +faststart final.mp4`.

2026 06 15

§ . . — Project-name-prefixed
feature/QA recording filenames (User mandate,
- -)

Forensic anchor — verbatim user mandate (2026-06-15): "All recorded videos MUST START with prefix: the PROJECT NAME (ALWAYS USE THE PROJECT NAME). Project name MUST be obtained according to the constitution's own project-name resolution."

Every recorded video the project produces — every feature/QA real-use recording (§11.4.153), every window-scoped capture (§11.4.154), every always-on device recording (§11.4.128), and every raw or curated recording artefact under the project-declared recording path (§11.4.35) and the committed

`docs/qa/<run-id>/` evidence trail (§11.4.83) — MUST have a filename that STARTS WITH the PROJECT-NAME prefix, ALWAYS, with no exception. A recording filename that omits the project-name prefix is a §11.4.155 violation: a corpus spanning multiple projects/scopes on one host (a real condition under §11.4.128 always-on recording + §11.4.103 parallel streams) becomes un-greppable and un-attributable, and a reader cannot tell at a glance which project a recording belongs to — the same identify-and-grep failure §11.4.151 forbids on the release-tag axis, applied to the recording-corpus axis.

Prefix resolution order (closed-set, deterministic — §11.4.6 no-guessing;

IDENTICAL to §11.4.151's prefix resolution): the project-name prefix MUST be resolved, never guessed, by the constitution's own project-name resolution: (1)

`HELIX_RELEASE_PREFIX` from the project's `.env` — authoritative when set; `.env` is git-ignored per §11.4.30 and the variable is documented in the tracked `.env.example` (a §11.4.77 re-obtain mechanism, never committed); (2) fallback = the lowercased snake_case form of the project root directory name (no spaces) per §11.4.29 — used whenever the env var is unset/empty, so a prefix is ALWAYS resolvable from the checkout with zero operator input. The SAME resolved prefix is used for EVERY recording the project produces in a given checkout, so a single `ls '<PREFIX>---'*` (or `find . -name '<PREFIX>---*'`) enumerates the whole recording corpus for that project. Canonical filename form: `<PREFIX>---<feature-or-scope>---<run-id>.<ext>` (the `---` triple-hyphen separator keeps the prefix unambiguously delimited from a feature/scope name that may itself contain hyphens). The prefix MUST be the SAME value §11.4.151 resolves for release tags in the same checkout — a recording prefix and a release-tag prefix diverging in one checkout is itself a §11.4.155 violation (one project, one resolved name).

Honest boundary (§11.4.6): the project-name prefix guarantees a recording is attributable + greppable to its project, NOT that the recording's CONTENT is valid — content validity still rests on the §11.4.107 liveness battery, the §11.4.137 content-correctness oracle, and the §11.4.153 video-analysis remediation loop; the prefix is a naming/attribution discipline that composes with, never replaces, those evidence layers. The prefix also does NOT relax §11.4.154's window-scoped-

capture + fresh-corpus-rotation invariants (rotation removes the agent's OWN prior in-scope `<PREFIX>---*` recordings first; a foreign-prefix or operator-authored file is surfaced, never deleted, per §11.4.122 + §9.2).

Classification: universal (§11.4.17) — a platform-neutral recording-attribution discipline reusable by ANY project that produces recordings; the consuming project supplies its concrete prefix value + the `HELIX_RELEASE_PREFIX` env var + its recording path per §11.4.35. Composes §11.4.151 (the SAME prefix-resolution order + the same identify-and-grep purpose, applied to recordings instead of release tags) / §11.4.128 (every always-on device recording filename carries the prefix) / §11.4.153 (the per-feature real-use video's confirmation path is prefixed) / §11.4.154 (window-scoped + fresh-corpus rotation operate on prefixed filenames) / §11.4.111 (the prefix is a stable-identity name, not an enumeration index) / §11.4.83 (committed `docs/qa/<run-id>/` recording evidence carries the prefix) plus §11.4.6 / §11.4.29 / §11.4.30 / §11.4.35 / §11.4.77 / §11.4.86. Propagation gate `CM-COVENANT-114-155-PROPAGATION` (literal `11.4.155` across the consumer fleet) + recommended gate `CM-RECORDING-PROJECT-NAME-PREFIX` (every recording filename at the project's recording path + every committed `docs/qa/<run-id>/` recording artefact starts with the resolved `<PREFIX>---` prefix, identical to the §11.4.151-resolved prefix for the checkout) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; write a recording with no project-name prefix, or a prefix differing from the §11.4.151-resolved value → `CM-RECORDING-PROJECT-NAME-PREFIX` FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.155. Non-compliance is a release blocker. No escape hatch — no `--no-recording-prefix`, `--recording-without-project-name`, `--unprefixed-recording`, `--prefix-optional-for-recording`, `--differing-recording-prefix` flag.

§ . . — All CI/CD automation (GitHub Actions / GitLab pipelines / equivalents) MUST be disabled (User mandate, — —)

Forensic anchor — verbatim user mandate (2026-06-15):

"Any GitHub actions or GitLab pipelines MUST BE disabled! Add this critical mandatory rule / mandatory constraint into the root constitution Submodule, commit and fetch all its changes to all upstreams and make sure we respect and follow this rule (we do apply it) ASAP!!!"

Every repository this Constitution governs — the main repo, this constitution submodule, and every owned + nested submodule we author and push — MUST ship with ALL server-side CI/CD automation DISABLED. No push to any owned upstream may trigger a GitHub Actions run, a GitLab pipeline, or any equivalent provider-side automation (Jenkins, CircleCI, Travis, Drone, Woodpecker, Bitbucket Pipelines, Azure Pipelines, or any `on: push / schedule / workflow_dispatch` workflow). This GENERALISES + makes ABSOLUTE the §11.4.75 Layer-5 posture ("Remote CI surfaces ... are DISABLED — the workflow file is preserved at ... `disabled-local-only` (NOT `.yaml` ; GitHub Actions ignores it)": what §11.4.75 did for the single constitution-compliance workflow, §11.4.156 mandates for ALL CI in ALL governed repositories. Enforcement migrates to the LOCAL §11.4.75 five-layer git-hook ritual + the §11.4.40 pre-tag sweep — never a remote runner.

The mandate (ALL hold): **(A) Zero active CI at the repository root.** No active `.github/workflows/*.yaml|*.yml`, no `.gitlab-ci.yml`, no `.gitlab/**` pipeline include, nor any equivalent provider config may exist at the ROOT of any governed repository/submodule — the only location a provider executes. **(B) Disabled means a push triggers ZERO runs.** Delete the config OR rename it to a non-active name (the §11.4.75 `.disabled` / `.disabled-local-only` convention — a provider ignores any name that is not its exact trigger filename); a workflow left with live `on: triggers` but `if: false` jobs still queues runs and is NOT compliant. **(C) Scope = repositories we author + push.** Vended / third-party source whose nested `.github/workflows` or `.gitlab-ci.yml` sit BELOW the repo root (e.g. AOSP `external/**`, `prebuilts/**`, `packages/modules/**`, vended submodules) are INERT — a provider never executes a non-root config — so they are OUT of scope and MUST NOT be mass-edited (§11.4.29 vendor-name exemption + tree-integrity); the test is "does a push to one of OUR upstreams trigger a run?" — yes ⇒ disable, structurally-inert ⇒ document + leave (§11.4.6 — verify inertness as FACT, never assume). **(D) No new CI may be added.** Introducing any active workflow/pipeline into a governed repo is a release blocker. **(E) Pre-push verification.** Before any push, verify no root-level active CI exists in

the pushed repo/submodule (`git ls-files | grep -E '^\.github/workflows/.*\.ya?ml$|^\.gitlab-ci\.yaml$'` returns empty for authored repos); a §11.4.109-class PreToolUse guard + the recommended gate enforce the class mechanically. Honest boundary (§11.4.6): file-level disabling stops FILE-triggered runs; it does NOT disable provider-side server settings the agent cannot reach (org-default required workflows, branch-protection required checks, provider-side scheduled exports) — those MUST be turned off in provider settings by the operator, and the agent documents what it cannot reach rather than claiming a completeness it did not achieve.

Classification: universal (§11.4.17) — a platform-neutral CI-governance discipline reusable by ANY project; the consuming project supplies its concrete repository/submodule set + provider list per §11.4.35. Composes §11.4.75 (mechanical-enforcement Layer 5 — §11.4.156 promotes its CI-disabled posture to an absolute universal rule across all governed repos) / §11.4.29 (vendor/third-party name + tree exemption) / §11.4.6 (verify inertness, never guess) / §11.4.40 / §11.4.42 (the LOCAL pre-tag ritual replaces the remote runner) / §11.4.109 (PreToolUse guard blocks the class at the tool-call boundary) / §11.4.113 (same no-remote-surprise discipline applied to pushes) / §2.1 (multi-upstream push — the disabled state holds on every mirror). Propagation gate `CM-COVENANT-114-156-PROPAGATION` (literal `11.4.156` across the consumer fleet) + recommended gate `CM-NO-ACTIVE-CI` (no root-level active CI config in any authored repo/submodule, verified pre-push) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; add a root `.github/workflows/x.yaml` to an authored repo → `CM-NO-ACTIVE-CI` FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.156.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--allow-ci` , `--enable-workflow` , `--keep-pipeline` , `--remote-ci-OK` , `--ci-exempt` flag.

2026 06 15

§ . . — GEMINI.md maintained in lockstep with CLAUDE.md / AGENTS.md / QWEN.md (User mandate, — —)

Forensic anchor — verbatim user mandate (2026-06-15):

"Make sure with CLAUDE.md, AGENTS.md, QWEN.md we maintain GEMINI.md too! Add this mandatory fact / rule to root constitution Submodule we are inheriting / extending - CONSTITUTION.md, CLAUDE.md, QWEN.md, AGENTS.md, GEMINI.md and other related relevant files!"

Forensic case study (FACT, 2026-06-15). When §11.4.156 + §11.4.157 were authored, Constitution.md / CLAUDE.md / AGENTS.md / QWEN.md carried the rule family through §11.4.155, but GEMINI.md had silently drifted to §11.4.141 — fourteen mandates (§11.4.142–§11.4.155) never propagated to it. A per-agent context carrier that lags is a §11.4 propagation-bluff: a Gemini-CLI agent reading the stale GEMINI.md operates under an out-of-date Constitution while the fleet believes the rule is universally in force.

GEMINI.md is a FIRST-CLASS governance context carrier, EQUAL to CLAUDE.md / AGENTS.md / QWEN.md — NEVER an optional or best-effort sibling. Every governance addition or edit MUST land in ALL FIVE carriers in lockstep: the canonical Constitution.md PLUS each per-agent mirror CLAUDE.md + AGENTS.md + QWEN.md + GEMINI.md (and any other relevant per-agent context file a runtime introduces), in the SAME change-window, each with its synchronized .html / .pdf / .docx exports per §11.4.65. The mandate (ALL hold): **(A) Five-carrier lockstep.** No governance change is complete until GEMINI.md carries it alongside the other three mirrors — GEMINI.md is added to the §11.4.26 constitution-update pipeline's stage-2 propagation set + stage-3 cross-reference validation explicitly. **(B) No silent drift.** A GEMINI.md whose highest rule number lags the other mirrors is a §11.4.157 violation (the same severity class as a §11.4.65 stale-export drift); the gap MUST be back-filled. **(C) Equal status, equal rules.** GEMINI.md restates the SAME literal 11.4.N anchors the §11.4.x propagation gates require (per §11.4.35 inheritance), so the propagation-gate fleet count INCLUDES GEMINI.md. **(D) Consumer projects too.** The lockstep binds the consuming project's repository-root context files (its own CLAUDE.md / AGENTS.md / QWEN.md / GEMINI.md) per §11.4.35 — a project that maintains the other three but not GEMINI.md is non-compliant. Honest boundary (§11.4.6): the pre-existing §11.4.142–§11.4.155 GEMINI.md back-fill is a tracked remediation

work unit; this rule MAKES that drift a release-blocking debt rather than silently tolerating it — claiming `GEMINI.md` is "in sync" while the back-fill is incomplete is itself a §11.4.157 violation.

Classification: universal (§11.4.17) — a platform-neutral context-carrier-parity discipline reusable by ANY project that ships per-agent context files. Composes §11.4.26 (constitution-update workflow — `GEMINI.md` added to the propagation + validation set) / §11.4.35 (canonical-root inheritance — the per-agent mirrors restate the canonical anchors) / §11.4.17 (every new rule lands in all carriers) / §11.4.44 / §11.4.65 (revision header + synchronized exports for `GEMINI.md`) / §11.4.140 (`GEMINI.md` is one of the "every CLI agent" context carriers) / §11.4.156 (this rule's own first lockstep test) / §1.1. Propagation gate `CM-COVENANT-114-157-PROPAGATION` (literal `11.4.157` across the consumer fleet, `GEMINI.md` INCLUDED) + recommended gate `CM-GEMINI-MD-LOCKSTEP` (`GEMINI.md`'s highest §11.4.N is not less than the other three mirrors' highest) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; let `GEMINI.md`'s highest rule number fall behind the other mirrors → `CM-GEMINI-MD-LOCKSTEP` FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.157.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--skip-gemini-md`, `--gemini-optional`, `--gemini-lag-OK`, `--four-carrier-suffices` flag.

2026 06 16

§ . . — Intensive all-feature/flow/
edge-case video-recording + read-the-screen
content-verification mandate (User mandate,
— —)

Forensic anchor — verbatim user mandate (2026-06-16): "Add the following mandatory rules / guidelines / mandatory constraints ... regarding the recording of all features, flows, cases, edge cases and full validation and verification of it all with real results, covered with evidence and rock solid physical proofs and no bluffs of any kind! All projects MUST BE covered with such intensive testing and video recordings produced as the results! Default save path of them all is ALWAYS to be Downloads directory under host machine's user's home directory! All video

recordings MUST contain valid working features with no false results showed in it and all printed out logs, messages, UI labels, dialogs, toasts MUST BE really read by the testing System (HelixQA MUST BE used!!!) so there are no issues or bluff!"

Forensic case study (FACT, 2026-06-16): the §11.4.153 feature-video sweep recorded `cli -stream` and the recording SHOWED `streaming generation failed: invalid character 'd' looking for beginning of value` — a REAL failure the sweep caught precisely because the on-screen text was READ + analyzed, not merely "a video was produced". Root cause was a §11.4.108 stale binary (source-fixed at commit 74077736, deployed artifact stale); the recording-that-reads-the-screen turned an otherwise-invisible artifact gap into a caught defect (then fixed + §11.4.135-guarded). This is the value §11.4.158 makes universal: a recording is proof ONLY when its shown content is machine-read and verified to be a genuine working result.

Strict intensification of §11.4.153 (per-feature Status+video ledger) + §11.4.107 (liveness + analyzer self-validation) + §11.4.117 (CV/OCR pixel oracle) + §11.4.137 (content-correctness oracle) + §11.4.27/§11.4.52 (HelixQA + autonomous validation), binding them into ONE intensive-coverage record-and-read mandate. ALL hold:

(A) Intensive total coverage. Every project MUST be covered by intensive automated testing that EXERCISES + RECORDS every feature, every flow, every use case, AND every edge case (valid / invalid / boundary / empty / max / concurrent / failure-injection per §11.4.85) — not a happy-path sample. A feature/flow/edge-case with no recorded, read-verified evidence is a §11.4 PASS-bluff at the coverage layer (composes §11.4.25 full-automation-coverage + §11.4.118 discovery-completeness + §11.4.153 per-feature ledger).

(B) Real results, rock-solid physical proof, zero bluff. Every recording MUST show the feature GENUINELY WORKING with REAL results (real prompts → real LLM/service responses → real outputs per §11.4.153) — NO false / simulated / stale / frozen result may appear in any recording (§11.4.2/.5/.107). A recording that shows an error, a bluff response, a frozen/stale frame, or a non-working feature is a FINDING (→ §11.4.153(4) / §11.4.4 fix → retest → re-record), NEVER a confirmation.

(C) Read-the-screen content verification (HelixQA-driven). The testing System MUST ACTUALLY READ the content shown in each recording — every printed log line, message, UI label, dialog, toast, status text — and VERIFY it (OCR/ROI per

§11.4.117 / §11.4.137(12) with a confidence floor for pixel surfaces; direct text capture for terminal/log surfaces; the §11.4.107(10) golden-good/golden-bad self-validated analyzer). "A video was produced" is NOT evidence; "the System read the on-screen text and confirmed it is a genuine working result" is. HelixQA (§11.4.27, the HelixDevelopment/HelixQA submodule) MUST be the driver/orchestrator of this exercise → record → read → score pass — its banks exercise the features, capture the recordings, read the shown content, and score PASS only on a read-confirmed genuine result with the captured artefact path (§11.4.69).

(D) Default save path = `$HOME/Downloads` . Unless a project declares an explicit override per §11.4.35, the default save path for ALL recordings is the host user's home Downloads directory (`$HOME/Downloads` — resolved from the running user's home at runtime, NEVER hardcoded). A project MAY override the location per §11.4.35 (e.g. an external-drive recordings root); the override is recorded in the project layer, not here. Filenames carry the §11.4.155 project-name prefix; window-scoped capture + fresh-corpus rotation per §11.4.154; raw corpus git-ignored + code-intelligence-excluded per §11.4.128, only curated evidence committed under `docs/qa/<run-id>/` (§11.4.83) at release prep.

(E) No false results survive. Composes the whole §11.4 anti-bluff covenant — a recording-confirmed PASS the operator (or a later re-read) contradicts is a §11.4.138 operator-escape → mandatory bluff-audit + permanent guard. The read-and-verify analyzer is itself anti-bluff (golden-good PASSES, golden-bad FAILS, §1.1).

(F) Vision analysis MANDATORY for every recording. After every recording, the agent MUST perform vision analysis: read the terminal output / screenshot frames / video content and verify (i) the feature ACTUALLY WORKS (real output, not errors, not empty), (ii) LLM responses are REAL (not simulated, not placeholder, not "for now"), (iii) all tests show PASS with captured evidence, (iv) no "TODO implement", "simulate", "for now", or "in production this would" patterns appear, and (v) the output demonstrates the feature working for the end user. Vision validation MUST produce a verdict (PASS/FAIL) with cited evidence path (§11.4.69). A recording without vision validation is a §11.4.158 violation — a video that was never read is NOT evidence.

Honest boundary (§11.4.6): reading + verifying the shown content proves the recorded run produced a genuine working result for the cases exercised — it does NOT prove zero remaining defects (§11.4.118) and does NOT replace §11.4.108 runtime-signature verification or §11.4.40 full-suite retest; it is the recorded-evidence layer every feature additionally crosses.

Classification: universal (§11.4.17) — the consuming project supplies its concrete capture mechanism, OCR/read tooling, HelixQA banks, and (optional) recording-path override per §11.4.35. Composes §11.4.2 / §11.4.5 / §11.4.25 / §11.4.27 / §11.4.52 / §11.4.69 / §11.4.83 / §11.4.85 / §11.4.107 / §11.4.108 / §11.4.117 / §11.4.118 / §11.4.128 / §11.4.137 / §11.4.138 / §11.4.153 / §11.4.154 / §11.4.155 / §1.1. Propagation gate CM-COVENANT-114-158-PROPAGATION (literal 11.4.158) + recommended gates CM-INTENSIVE-RECORDING-COVERAGE + CM-RECORDING-CONTENT-READ-VERIFIED + paired §1.1 meta-test mutation.

Canonical authority: constitution submodule Constitution.md §11.4.158. Non-compliance is a release blocker. No escape hatch — no --skip-recording-coverage , --video-without-content-read , --happy-path-recording-suffices , --recording-path-anywhere , --unread-recording-OK , --skip-helixqa-read flag.

2026 06 20

§ . . — Mandatory window-specific video recording + vision validation mandate (User mandate, - -)

Forensic anchor — verbatim user mandate (2026-06-20): "All video evidence must be window-specific MP4 with mandatory vision validation, terminal cleanup, real results only."

Every feature test, validation, verification, challenge, and QA session that produces video evidence MUST comply with ALL of the following:

(A) Window-specific recording ONLY. Every video MUST record ONLY the target application window (Terminal pane, TUI app, browser tab, emulator frame), NEVER the whole desktop/monitor screen. Use window-specific capture mechanisms: macOS screencapture -l<window_id> (obtain window id via osascript -e 'tell application "System Events" to get id of first window of process "<app>"'), Linux xdotool + ffmpeg -video_size WxH -f x11grab -i :0.0+X,Y , or equivalent platform-specific window-targeted capture. Whole-desktop

capture leaks operator-private content (§11.4.10), dilutes the §11.4.107 liveness/freeze oracle, and breaks the §11.4.137 OCR/ROI content oracle. Platform genuinely cannot capture below whole-screen => honest §11.4.3 SKIP + tracked migration item, never a whole-screen pass-off.

(B) MP4 format REQUIRED. Every recording MUST be saved as `.mp4` format (H.264 codec, `movflags +faststart`, `pix_fmt yuv420p`). `.cast` (asciinema) files are supplementary only — the primary evidence is the `.mp4` video. If using asciinema for terminal recording, auto-convert to `.mp4` via `agg + ffmpeg` immediately after capture per §11.4.154(C).

(C) Project-name prefix REQUIRED. Every recording filename MUST start with the project name in snake_case, derived from `HELIX_RELEASE_PREFIX` in `.env` (per §11.4.151) or the lowercased project root directory name (per §11.4.29). Format: `<project_name>-<feature>-<YYYYMMDD-HHMMSS>.mp4`. An unprefixed recording is a §11.4.159 violation (same severity class as §11.4.155).

(D) Mandatory vision validation. After EVERY recording, the agent MUST perform vision analysis: read the terminal output / screenshot frames / video content and verify: (i) the feature ACTUALLY WORKS (real output, not errors, not empty); (ii) LLM responses are REAL (not simulated, not placeholder, not "for now"); (iii) all tests show PASS with captured evidence; (iv) no "TODO implement", "simulate", "for now", or "in production this would" patterns appear; (v) the output demonstrates the feature working for the end user. Vision validation MUST produce a verdict (PASS/FAIL) with cited evidence path (§11.4.69). A recording without vision validation is a §11.4.159 violation — a video that was never read is NOT evidence.

(E) Terminal window cleanup. After each recording completes, the agent MUST dismiss/close ONLY the Terminal window used for that recording (using window-specific close: `osascript` with window id, `xdotool`, or equivalent). MUST NOT close Terminal windows belonging to other working processes. A recording that leaves orphan Terminal windows is a §11.4.159 violation.

(F) Real results ONLY. Every recording MUST show REAL working features — real API calls, real LLM responses, real test results. A recording showing errors, empty output, simulated responses, or placeholder content is a §11.4.159 violation and MUST trigger a fix -> retest -> re-record cycle before acceptance.

(G) Re-runnable evidence. Every recording MUST be reproducible: the command shown in the video MUST be re-runnable to produce the same results. Recordings of one-time manual interactions that cannot be automated are §11.4.159 violations.

(H) Fresh-corpus rotation. When a new recording run for a scope begins, the agent's own prior in-scope stale recordings at the recording path MUST be removed FIRST (per §11.4.154). Committed `docs/qa/<run-id>/` evidence is the durable record, NOT rotated.

(I) Content verification MANDATORY — not duration-based. The value of a recording is NOT its duration but its CONTENT. A recording MUST demonstrate the ACTUAL feature being used with REAL results. Before accepting ANY recording, the agent MUST verify:

- The expected output patterns ARE present in the recording (e.g., test PASS lines, API response data, feature-specific output)
- The feature ACTUALLY WORKS as demonstrated (not just "something ran")
- LLM responses are REAL content (not simulated, not placeholder, not empty)
- Every claim of "working" is backed by visible evidence in the recording

A 5-second recording that shows a feature working correctly is MORE valuable than a 60-second recording of empty terminal. Duration is NOT a proxy for quality.

(J) Expected-content specification REQUIRED. Before recording, the agent MUST specify what content SHOULD appear in the recording (expected patterns, expected test results, expected API responses). After recording, the agent MUST verify these patterns ARE present. If expected content is MISSING, the recording is REJECTED regardless of duration.

(K) Content-verification recording workflow. The mandated workflow for every recording:

1. SPECIFY expected content patterns (what SHOULD appear)
2. RECORD the feature execution
3. EXTRACT all text from the recording
4. VERIFY expected patterns are present
5. CHECK for simulated/placeholder content
6. ACCEPT only if ALL patterns found AND zero bluffs detected

7. REJECT and re-record if ANY pattern missing or bluff detected

(L) Root cause analysis REQUIRED for rejected recordings. When a recording is rejected (missing expected content, bluff detected, or empty capture), the agent MUST investigate WHY before re-recording. Per §11.4.102, determine the root cause (timing issue, wrong command, tool failure, etc.) and fix it. Simply re-recording without understanding WHY the first attempt failed is a §11.4.159 violation.

(M) Real-time monitoring RECOMMENDED. For complex features, use real-time monitoring that analyzes output DURING recording (not after). This catches issues immediately and allows corrective action before the recording completes.

Classification: universal (§11.4.17) — the consuming project supplies its concrete window-capture mechanism (per surface class), recording path, project-name prefix, and vision-validation tooling per §11.4.35. Composes §11.4.2 / §11.4.3 / §11.4.5 / §11.4.10 / §11.4.29 / §11.4.69 / §11.4.83 / §11.4.107 / §11.4.111 / §11.4.128 / §11.4.137 / §11.4.151 / §11.4.153 / §11.4.154 / §11.4.155 / §11.4.158 / §1.1. Propagation gate CM-COVENANT-114-159-PROPAGATION (literal 11.4.159) + recommended gate CM-WINDOW-VIDEO-VALIDATED (every video is window-specific MP4 with vision-confirmed verdict + terminal cleanup + project-prefix) + paired §1.1 meta-test mutation (strip the literal -> propagation gate FAILs; a whole-screen recording where window-scoped is feasible, or a recording without vision validation -> CM-WINDOW-VIDEO-VALIDATED FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md §11.4.159. Non-compliance is a release blocker. No escape hatch — no --whole-screen-ok , --cast-only , --skip-vision-validation , --no-cleanup , --simulated-recording-ok , --unprefixed-recording flag.

§ . . — Vision-verified recording + HelixQA bridge mandate (User mandate, — —)

Forensic anchor — direct user mandate (2026-06-21). Every video recording produced for feature/QA evidence under this Constitution MUST be processed through a vision/OCR pipeline that reads the on-screen content and confirms expected results BEFORE the recording is accepted as evidence. A recording that

is accepted without automated read-the-screen verification is a §11.4 PASS-bluff at the evidence-integrity layer — "a video was produced" does not prove the video contains a genuine working result, and operator-side manual review does not scale to the intensive recording mandate (§11.4.158).

The bridge requirement. The recording system MUST provide a real-time or post-recording bridge that feeds captured frames/pixels to HelixQA's test infrastructure (or equivalent automated test framework) for content verification — automated read-the-screen verification against expected content patterns specified BEFORE recording (per §11.4.159(J)). The bridge MUST:

1. **Capture frames at ≤5 s intervals** during the recording window — sufficient to assert the feature's progression, not merely its start and end state.
2. **Run OCR/vision analysis against each frame** using a self-validated golden-good/golden-bad analyzer per §11.4.107(10) — the analyzer itself MUST be mutation-tested so it cannot bluff a PASS on empty/degraded content.
3. **Compare extracted text against the SPECIFY-phase expected patterns** — every pattern from the §11.4.159(J) expected-content specification MUST be checked.
4. **Produce a per-frame PASS/FAIL verdict with an evidence path to the frame** — raw frame image + OCR output + comparison result, linked from the verdict.
5. **Surface failures immediately** so the recording can be re-done per §11.4.159(L) root-cause analysis, not queued for batch review after the recording window closes.

Project-configured parameters (per §11.4.35). The frame-capture interval, OCR confidence floor, and pattern-matching thresholds are project-configured, calibrated on the project's own fixtures per §11.4.6 (no hardcoded industry defaults).

Honest boundary (§11.4.6). Vision verification confirms the feature produced expected on-screen output — it does NOT replace §11.4.5 captured-evidence quality analysis (audio RMS/XRUN, video freeze/jitter/obstruction census) nor §11.4.108 runtime-signature verification (the artifact's bytes landed and are active on a clean target). If the recording surface is FLAG_SECURE / DRM-protected / sink-blanked per §11.4.112, the test documents the gap and uses the §11.4.117 proxy oracle (accessibility cue / subdebug channel) instead — never a faked pixel-verified PASS.

Composition. §11.4.160 composes with §11.4.5 (captured-evidence quality — §11.4.160 adds the automated read-the-screen layer, does not replace quality metrics), §11.4.27 (HelixQA is the framework that runs the bridge), §11.4.69 (verdict with evidence path), §11.4.107(10) (self-validated analyzer), §11.4.112 (secure-surface gap documented), §11.4.117 (proxy oracle for non-introspectable surfaces), §11.4.153 (per-feature video ledger), §11.4.158 (intensive recording coverage — §11.4.160 adds the verification bridge to every recording), §11.4.159(J) (SPECIFY-phase expected-content patterns are the bridge's input).

Gates. Propagation gate `CM-COVENANT-114-160-PROPAGATION` (literal `11.4.160`) + recommended gate `CM-VISION-VERIFIED-RECORDING-BRIDGE` (every recording has a self-validated OCR/vision pass against SPECIFY-phase patterns, with per-frame PASS/FAIL) + paired §1.1 meta-test mutation (strip the bridge/accept a recording without vision verification → gate FAILs; gate-code = separate work item).

Classification: universal (§11.4.17) — the consuming project supplies its concrete capture mechanism, OCR/vision tooling, capture interval, confidence floor, and HelixQA bridge integration per §11.4.35. No escape hatch — no `--skip-vision-verification`, `--manual-review-suffices`, `--no-bridge`, `--single-frame-suffices`, `--hardcoded-thresholds-OK`, `--unvalidated-analyzer-OK` flag exists.

§ . . — Rootless container runtime
mandate (User mandate, - -)

Forensic anchor — direct user mandate (2026-06-21).

The mandate. Every project governed by this Constitution MUST use **Podman in rootless mode** (or an equivalent rootless container runtime) for ALL containerized workloads. Docker in rootful mode, `sudo`, or any escalation to root for container management is **STRICTLY FORBIDDEN** unless the target platform genuinely has no rootless container option AND that platform constraint is documented per §11.4.112 (structurally-impossible — with cited authoritative sources proving the absence of a rootless option on that platform). Rootless container execution eliminates the most common container-to-host privilege-escalation vector, aligns

with the security-by-default posture of the §11.4.10 credentials-handling mandate, and is available and production-mature on Linux (Podman rootless since v2.0+, 2020) and macOS (Podman Machine with `--rootful=false`).

The `vasic-digital/containers` Submodule as the sole orchestration layer.

Per §11.4.76 (Containers-submodule mandate), the `vasic-digital/containers` Submodule MUST be used as the **sole container orchestration layer** — every container operation MUST go through the Submodule's `pkg/boot` / `pkg/compose` / `pkg/health` primitives. No ad-hoc `docker` / `podman` commands outside the Submodule; if the Submodule's `pkg/boot` does not support a lifecycle or runtime the project needs, the missing capability MUST be extended in the Submodule itself via upstream PR (the §11.4.74 extend-don't-reimplement discipline), not worked around with a raw container command in the project.

The on-demand-infra invariant (refines §11.4.76). All container-related integration tests MUST boot infra on-demand via the Submodule (`pkg/boot` is the test entrypoint), not depend on a pre-running container daemon or operator-side manual `podman machine start` . Operators are NEVER required to start `podman machine` / `docker compose up` / `podman-compose up` manually — the boot is part of the test entry point. A test that claims to exercise containerized components but silently depends on an already-running daemon (skipping the boot) is a §11.4 PASS-bluff at the test-infrastructure layer.

Honest boundary (§11.4.6). Rootless container execution eliminates the container-to-root privilege-escalation vector and enforces containerization hygiene — it does NOT replace §11.4.10 credentials-handling (credentials stored in containers are still subject to the leak audit and `.gitignore` mandates) nor §12.3 container hygiene (memory limits, OOM policy, restart backoff still apply). A platform that genuinely lacks rootless containers is documented per §11.4.112 and uses the Submodule's rootful mode with a documented extra risk acceptance — never uses ad-hoc rootful Docker.

Composition. §11.4.161 composes with §11.4.76 (Containers-submodule — §11.4.161 adds rootless-as-default + on-demand-infra invariant), §11.4.74 (extend-don't-reimplement — missing Submodule capabilities are extended upstream), §11.4.10 (credentials — rootless containers still require credential hygiene), §11.4.112 (structurally-impossible — the documentation path for platforms with no

rootless option), §12.3 (container hygiene — memory limits/OOM/restart policies still apply), §11.4.6 (no-guessing — "the target platform probably has rootless" is not a finding).

Gates. Propagation gate `CM-COVENANT-114-161-PROPAGATION` (literal `11.4.161`) + recommended gate `CM-ROOTLESS-CONTAINER-RUNTIME` (every container operation is verified as Podman rootless OR has a §11.4.112 struct-impossible citation for that platform; no `sudo /rootful Docker` for container management; container integration tests boot on-demand via the `containers` Submodule) + paired §1.1 meta-test mutation (replace a rootless invocation with `sudo` or ad-hoc docker, OR bypass the Submodule's `pkg/boot` with a raw `podman` command → gate FAILs; gate-code = separate work item).

Classification: universal (§11.4.17) — the consuming project supplies its concrete rootless runtime (Podman rootless or equivalent), the `containers` Submodule integration, and any platform constraint documentation per §11.4.112 per §11.4.35. No escape hatch — no `--allow-rootful`, `--sudo-container-OK`, `--ad-hoc-docker-permitted`, `--no-containers-submodule`, `--pre-started-daemon-OK` flag exists.

§ 11.4.161 — OpenDesign UI design system mandate (User mandate, 2026-06-21)

Forensic anchor — direct user mandate (2026-06-21).

The mandate. Every project governed by this Constitution that produces user-facing interfaces (web, desktop, mobile, TUI, or any other surface classes the project supports per §11.4.25) MUST use **OpenDesign** (<https://github.com/nexu-io/open-design>) as the mandatory UI design-and-refinement system — NOT ad-hoc CSS, not one-off design tools, not improvised inline styles, not a different design-token system. OpenDesign is consumed as a project dependency (npm package or equivalent per the consuming project's technology stack), never vendored or forked.

Design tokens and theming. The project MUST use OpenDesign's design tokens / themes system for:

1. **All color palette definitions** — supporting BOTH light and dark themes using the project's brand colors extracted from canonical brand assets (logo, style guides, brand guidelines per §11.4.35). Color tokens are defined in OpenDesign's token format, never duplicated in CSS custom properties or inline styles.
2. **Typography scale and spacing** — typefaces, font sizes, line heights, letter spacing, and content margin/padding tokens are defined in OpenDesign, consumed by every UI component, never scattered across individual component files.
3. **Component-level design tokens** — every reusable UI component (buttons, cards, dialogs, inputs, navigation elements, tables, modals, etc.) defines its tokens (border radius, shadow, hover/active/focus states, animation timing) in OpenDesign, not in component-specific stylesheets.

Extend, never reimplement. If a desired UI pattern (a component, theme variant, token type, or rendering target) is not supported by OpenDesign, the project MUST extend OpenDesign upstream per §11.4.74 (extend-don't-reimplement) — submit a PR adding the missing capability rather than working around it with ad-hoc CSS or a different tool. A pattern worked around outside the design system is a §11.4.162 violation (it silently fragments the UI surface — the divergence grows over time and the shared design system no longer governs all UI output).

Light+dark theme variants required. Every UI component MUST ship with BOTH light and dark theme variants, with the color palette defined from the project's canonical brand assets (logo/brand assets per §11.4.35). A component that renders correctly only in one theme is a §11.4.162 violation — a project's surfaces MUST be usable in both modes on any platform that supports theme switching.

Layout regression prohibition. UI elements MUST NOT overlap, fonts MUST NOT collide with other UI elements or other text, labels MUST NOT overlay labels. Any layout regression — including but not limited to clipped text, overlapping controls, misaligned grids, overflow-hidden content, or incorrect z-ordering — is a §11.4.162 violation regardless of whether the underlying token change had good

intent. The OpenDesign token-diff comparison is the primary regression-detection tool: token changes that produce layout regressions are caught by diffing the token set before and after the change, correlated with the project's visual regression tests.

Test coverage. All UI changes MUST be covered by the project's standard test types including visual regression tests (§11.4.4(b) layer 3) that capture before/after screenshots of every changed component/token, with per-pixel or perceptual-diff PASS/FAIL.

Honest boundary (§11.4.6). OpenDesign governs design tokens and theme definitions — it does NOT replace the project's functional testing (§11.4.27), accessibility testing (WCAG color contrast is still verified per §11.4.107), or the §11.4.48/§11.4.49 UI-driven and dual-approach testing methodology. A project that uses OpenDesign tokens but ships broken UI components is a §11.4.162 violation — the tool is mandatory, but using it correctly (tokens → components → tests → regressions) is equally mandatory.

Composition. §11.4.162 composes with §11.4.74 (extend-don't-reimplement — the rule for extending OpenDesign upstream), §11.4.25 (full-automation-coverage — UI components covered by every test type), §11.4.27 (no-fakes-beyond-unit — visual regression tests are real assertions, not mock screenshots), §11.4.4(b) (four-layer coverage includes visual regression at layer 3), §11.4.107 (accessibility/color-contrast assertions still apply), §11.4.48/§11.4.49 (UI-driven/dual-approach testing still applies), §11.4.35 (brand assets location and token-per-project instantiation), §11.4.6 (no-guessing — token configuration is declared, not inferred from appearance), §11.4.69 (visual regression evidence is captured per component).

Gates. Propagation gate `CM-COVENANT-114-162-PROPAGATION` (literal `11.4.162`) + recommended gate `CM-OPENDSIGN-UI-SYSTEM` (every UI-capable project has OpenDesign installed as a dependency, uses its token/theme system for all colors (light+dark), typography, spacing, and component tokens; extends OpenDesign upstream for unsupported patterns; no ad-hoc CSS-only UI; visual regression tests cover every UI change) + paired §1.1 meta-test mutation (replace a token-driven style with inline CSS → gate FAILs; ship a light-only component → gate FAILs; use a non-OpenDesign color tool → gate FAILs; gate-code = separate work item).

Classification: universal (§11.4.17) — the consuming project supplies its concrete OpenDesign package installation path, brand-asset location, token definitions file, component registry, and visual regression test harness per §11.4.35. No escape hatch — no `--skip-opendesign`, `--ad-hoc-css-OK`, `--light-only-OK`, `--no-token-system`, `--no-visual-regression-tests` flag exists.

2026 06 21

§ 11.4.17 — Universal Media Validation & Verification Mandate (User mandate, Forensic anchor)

Forensic anchor — direct user mandate (2026-06-21).

The mandate. Every recorded artifact (video MP4, audio WAV, screenshots PNG, asciinema cast, text output) produced by any project governed by this Constitution MUST pass through a MEDIA VALIDATION pipeline before being accepted as evidence. A recording that has NOT been validated by the pipeline is a §11.4 PASS-bluff at the evidence layer — it asserts the content is genuine and correct without having verified it.

(a) Content extraction. The pipeline MUST extract content from the media — OCR for video/screenshots (per §11.4.117 with a confidence floor per §11.4.107(12)), transcription for audio, text parsing for asciicast and text output. The extraction tooling is self-validated per §11.4.107(10) — a golden-good fixture MUST produce a PASS, a golden-bad fixture MUST produce a FAIL; an extractor that passes its golden-bad fixture is a bluff gate.

(b) Pattern matching. The extracted content MUST be compared against the SPECIFY-phase expected patterns per §11.4.159(J). Every specified pattern must be matched by at least one extracted line/frame; any specified pattern with no match yields a FAIL with an exact pinpoint record (which pattern, what the extracted content contained instead). The pattern set is the closed set of expected strings the test author declares before recording; an un-patterned PASS is inadmissible.

(c) Self-validated analyzer. The media validation pipeline as a whole (extraction + matching + verdict) MUST be mutation-tested with a golden-good/golden-bad fixture pair (§11.4.107(10)) — the golden-good MUST PASS, the golden-bad MUST FAIL. A pipeline that passes its golden-bad fixture is itself a §11.4 bluff gate; the pipeline MUST NOT be used for any evidence until its self-validation is restored.

(d) Structured verdict. The pipeline MUST produce a structured verdict for each artifact containing: PASS/FAIL, exact evidence path (file + frame/line/timestamp), matched patterns list (pattern → extracted text → frame/line), unmatched patterns list (pattern → expected → actual), analyzer version and calibration hash. The verdict is committed alongside the artifact under `docs/qa/<run-id>/` (§11.4.83).

(e) Exact pinpoint on FAIL. When the verdict is FAIL, the pipeline MUST provide pinpoint data — exactly which media artifact, which frame (video) or line (text/asciicast) or timestamp (audio), which OCR region or audio segment, which expected pattern, and what the extracted content contained instead. The pinpoint is the root-cause entry point for §11.4.102 systematic-debugging before re-recording per §11.4.159(L).

(f) Dual trigger mode. The pipeline MUST be triggerable both as a post-recording step (validate after the full recording completes, per §11.4.159(K)) AND as a real-time monitoring step that analyses frames/output during recording per §11.4.159(M) — real-time mode catches failures immediately, reducing re-record overhead.

(g) Paired §1.1 meta-test mutation. The media validator is itself subject to paired §1.1 mutation — a golden-bad fixture that the validator MUST reject, or the validator is a bluff gate. The mutation test creates a golden-bad fixture, feeds it to the validator, and asserts FAIL. If the validator returns PASS on the golden-bad fixture, the mutation test FAILs and the validator is inadmissible until fixed.

Honest boundary (§11.4.6). The media validation pipeline confirms the artifact contains the content it claims to contain and matches the intended patterns — it does NOT prove the feature is correct for the end user (that rests on §11.4.5 captured-evidence quality analysis, §11.4.107 liveness, and §11.4.108 runtime-signature), and it does NOT replace §11.4.40 full-suite retest. A pipeline that can be bypassed (a recording accepted without validation) is a §11.4.163 violation — the pipeline is mandatory for every recorded artifact, never optional.

Classification: universal (§11.4.17) — the consuming project supplies its concrete media extractors (OCR engine, transcription engine, text parser), golden-good/golden-bad fixture pair, expected-pattern declaration mechanism, and structured-verdict template per §11.4.35. Composes

§11.4.107(10)/.112/.117/.159(J)/.159(K)/.159(M)/.83/.5/.69/.102/.134/§1.1.

Propagation gate `CM-COVENANT-114-163-PROPAGATION` (literal `11.4.163`) +

recommended gate `CM-MEDIA-VALIDATION-PIPELINE` (every recorded artifact

crosses a self-validated media validation pipeline before acceptance; golden-bad fixture produces FAIL) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; bypass the pipeline and accept a raw unvalidated recording → CM-MEDIA-VALIDATION-PIPELINE FAILs; gate-code = separate work item).

2026 06 21

§ . . – Universal Constitution Auto-Propagation & Hook System (User mandate, – –)

Forensic anchor — direct user mandate (2026-06-21).

The mandate. Every fetch+pull of the constitution submodule MUST trigger an automatic post-update hook that installs/registers newly added or modified constitution components, so a pull of the constitution submodule is a COMPLETE governance-update transaction — not a partial one that leaves skills unregistered, MCP servers unknown, hooks unwired, or scripts non-executable. The canonical hook script lives at `constitution/scripts/post_update_hook.sh` and is inherited by reference (§11.4.28) — NEVER copied locally.

(a) Change detection. The hook MUST detect which constitution submodule files changed in the latest pull: `Constitution.md` , `CLAUDE.md` , `AGENTS.md` , `QWEN.md` , `GEMINI.md` , and any file under `scripts/` , `hooks/` , `skills/` , `mcp/` , `plugins/` . It compares HEAD before vs after the pull via `git diff --name-only HEAD@{1} HEAD` (or the equivalent for the merge/pull operation).

(b) Skill registration. For any newly added or modified file under `skills/` , the hook MUST install/register it with the CLI agent's skill system (e.g. for Claude Code: add it to the project's skill configuration so the agent discovers and can invoke it per §11.4.102(B)). The registration mechanism is project-runtime-specific per §11.4.35 — the hook invokes a consumer-defined script at a canonical path (`scripts/register_skills.sh`) that the consumer project provides.

(c) MCP server registration. For any newly added or modified file under `mcp/` , the hook MUST register the MCP server with the agent's MCP configuration (e.g. for Claude Code: add an entry to `.mcp.json` or `settings.json`). The registration mechanism is project-runtime-specific per §11.4.35 — the hook invokes a consumer-defined script at a canonical path (`scripts/register_mcp.sh`).

(d) Hook installation. For any newly added or modified file under `hooks/`, the hook MUST install it into the project's `.git/hooks/` or runtime hooks directory, making it executable. The hook script sets executable permission and copies/symlinks it; existing hooks with the same name are overwritten only if the new file's mtime is newer (safe default per §11.4.122 — never overwrite a consumer's custom hook without a newer upstream version).

(e) Script validation. For any newly added or modified file under `scripts/`, the hook MUST (i) make it executable, (ii) validate its syntax — `sh -n` for POSIX-sh scripts, `bash -n` for bash scripts (per §11.4.67), (iii) validate any `helix-deps.yaml` or `.gitignore-meta/` entries per §11.4.31/§11.4.77.

(f) Summary report. The hook MUST emit a summary of what was installed/updated: number of skills registered, MCP servers registered, hooks installed, scripts made executable and validated. The summary goes to stdout AND to `.constitution/state/last_update.log` (gitignored per §11.4.30).

(g) Graceful failure. If any component installation fails (skill not registered, MCP server not wired, hook not installed, script not validated), the hook MUST log the failure per §11.4.6 with exact file:line and continue with remaining installations — one failure does not abort the entire post-update transaction. The failure is surfaced in the summary report and is a finding for §11.4.102 investigation.

(h) Anti-bluff. The hook is itself anti-bluff — tested by a paired §1.1 mutation that adds/removes a skill in the constitution submodule, pulls the submodule, and asserts the hook detects the change. A hook that reports "nothing to install" when a new skill was added is a §11.4 PASS-bluff at the auto-propagation layer.

Consumer obligation. Every project that fetches/pulls the constitution submodule MUST invoke the post-update hook after the fetch completes. The invocation mechanism is project-defined per §11.4.35 — either a `post-update` git hook wired via `scripts/install_git_hooks.sh`, or a post-pull script that calls `constitution/scripts/post_update_hook.sh` explicitly. A project that pulls the constitution and does NOT invoke the hook is non-compliant — it has an incomplete governance-update transaction, and skills/MCP/hooks/scripts it believes are present may not be.

Honest boundary (§11.4.6). The auto-propagation hook installs/registers constitution components — it does NOT guarantee the consumer's runtime accepts them (MCP server may have incompatible dependencies, skill may require a

runtime version the consumer lacks) and does NOT replace §11.4.32 post-constitution-pull validation sweep (which verifies rules are enforced, not just that components are installed). Installation failures not investigated are the finding §11.4.102 captures.

Classification: universal (§11.4.17) — the consuming project supplies its concrete skill-registration script, MCP-registration script, hook-installation mechanism, and post-pull invocation method per §11.4.35. Composes

§11.4.28/.32/.35/.67/.77/.102(B)/.122/.31/§1.1. Propagation gate CM-COVENANT-114-164-PROPAGATION (literal 11.4.164) + recommended gate CM-CONSTITUTION-AUTO-PROPAGATION (post_update_hook.sh exists + executable, invoked after every pull, logs changes or failures, paired §1.1 mutation detects a new skill) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; pull the constitution and skip invoking the hook → CM-CONSTITUTION-AUTO-PROPAGATION FAILs; gate-code = separate work item).

2026 06 21

§ . . — Universal Independent Verification Agent Mandate (User mandate, - -)

Forensic anchor — direct user mandate (2026-06-21).

The mandate. Every code change OR recorded media artifact produced by any project governed by this Constitution MUST pass through an INDEPENDENT verification agent structurally separate from the author (§11.4.70/§11.4.20, dedicated subagent or distinct human, NEVER the author re-reading their own work — §11.4.92 self-evaluation PRECEDES and never satisfies). The verifier is one of MULTIPLE STRONG LAYERS — the FIRST verification layer every output crosses, complementing (never replacing) §11.4.108 runtime-signature verification, §11.4.125/§11.4.142 code review, and §11.4.40 full-suite retest.

(a) For CODE changes. The independent verification agent MUST: (i) run the §11.4.142 code-review discipline (adversarial analysis, blast radius, contract integrity, test genuineness), (ii) build and test the changed packages (invoke the project's pre_build_verification.sh on the diff scope), (iii) run the paired §1.1 meta-test mutations for every gate the change touches (asserting each mutated gate still FAILs and its restoration still PASSes), (iv) verify the §11.4.108 runtime-

signature registry entry for this change declares one machine-checkable observable. A code change that passes only one of (i)–(iv) is NOT verified — all four must hold.

(b) For MEDIA artifacts (video/audio/screenshots/asciicast/text output). The independent verification agent MUST: (i) run the §11.4.163 media validation pipeline on every recorded artifact, (ii) confirm the content is genuine — real LLM/service responses, not mocked/stub/simulated/placeholder (per §11.4.159(I)), (iii) confirm the artifact format is correct (MP4 H.264 for video, WAV/FLAC for audio, PNG for screenshots, cast+mp4 for asciicast per §11.4.159(B)), (iv) cross-check the evidence path against the artifact (the path exists, is non-empty, and the content matches the feature's expected output). A media artifact that passes format but shows simulated/placeholder content is a finding, not a PASS.

(c) For DOCS changes. The independent verification agent MUST: (i) validate that HTML + PDF + DOCX exports exist and are current per §11.4.65 (mtime parity), (ii) verify the §11.4.44 revision header is present and its integer has incremented from the previous commit, (iii) for docs that carry a §11.4.86 drift-proof fingerprint, verify the fingerprint matches the sorted roster/keyset. A doc change with a stale export or missing revision header is NOT verified.

(d) For CONFIG changes. The independent verification agent MUST: (i) validate YAML/JSON/TOML syntax (parse with the language's own parser — `yq` , `jq` , basic JSON parser, `python -m json.tool` , etc.), (ii) cross-reference against a schema where available (e.g. JSON Schema for `.settings.json` , YAML Schema for `helix-deps.yaml`), (iii) confirm no credential or secret pattern leaked per §11.4.10. A config change with a syntax error, schema violation, or leak is a finding, not a PASS.

(e) Structured findings. The verifier MUST produce a structured findings report — for each finding: item (which code/media/doc/config), verification class (a–d), tool/command used, evidence path (log file, build output, mutation result, validator verdict), PASS/FAIL. "Looks good" without evidence is not a finding — a rubber-stamp verifier is a §11.4 PASS-bluff at the verification layer.

(f) Iterate to GO. If any finding exists (PASS/FAIL indeterminate, BLOCKING, or even a warning), the verifier iterates per §11.4.134 — the author fixes the finding, the verifier RE-VERIFIES, until GO with ZERO findings and ZERO warnings. A single-pass verifier that reports findings but accepts GO because "the findings were minor" is a §11.4.165 violation.

(g) Anti-bluff — verifier self-validation. The independent verification agent is itself anti-bluff. The verifier's meta-test is a paired §1.1 mutation that introduces a deliberately broken code change (e.g. a gate with a stripped assertion) and asserts the verifier catches it. A verifier that passes a known-broken change is a bluff gate and must be investigated per §11.4.102 before any further use.

Honest boundary (§11.4.6). The independent verification agent confirms code/media/docs/config integrity at the source/artifact layer — it does NOT prove the feature works for the end user on the deployed target (that rests on §11.4.108 runtime-signature + §11.4.130 validate-fix-first + §11.4.40 full-suite retest) and does NOT replace the §11.4.40 full-suite retest or §11.4.118 discovery pass. A verifier that reports GO on an artifact whose downstream deployment is broken is not a verifier failure — that is caught by the later §11.4.108 layer. A verifier that reports GO without having actually run any tool (rubber-stamp) IS a verifier failure and a §11.4.165 violation.

Classification: universal (§11.4.17) — the consuming project supplies its concrete build+test command per diff scope, mutation-test runner, media validator invocation, doc export checker, config parser, and verifier-dispatch agent per §11.4.35. Composes

§11.4.142/.163/.44/.65/.86/.108/.125/.134/.10/.40/.92/.70/.20/.102/.159(I)/.159(B)/
§1.1. Propagation gate CM-COVENANT-114-165-PROPAGATION (literal 11.4.165)
+ recommended gate CM-INDEPENDENT-VERIFICATION-AGENT (every output — code/media/docs/config — passes an independent verifier structurally separate from the author, produces structured findings with evidence paths, iterates to zero-finding GO; verifier is self-validated by a paired §1.1 mutation) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; accept an output without the independent verifier's GO → CM-INDEPENDENT-VERIFICATION-AGENT FAILs; verifier reports GO on a known-broken fixture → self-validation FAILs; gate-code = separate work item).

§11.4.166 — REPEALED (operator decision, 2026-06-22). The Universal Semgrep static-analysis mandate is repealed: Semgrep is **NO LONGER a mandatory requirement** for projects governed by this Constitution. Static analysis / security scanning remains *encouraged* and may be re-adopted per project, but is NOT mandated, NOT a release blocker, and NOT cascaded. Removed with this repeal: `scripts/semgrep/` , `scripts/hooks/semgrep_precommit.sh` , the `submodules/semgrep` submodule, the MCP / pre-commit / PATH wiring, the consumer `.docs_chain/contexts/semgrep_status.yaml` context, and the `CM-COVENANT-114-166-PROPAGATION / CM-SEMGREP-WIRED` gates. Anchor number 11.4.166 is RETIRED (not reused).
^{11 4 167} Repealed before full fleet cascade — consuming projects never received it as a binding mandate.

§ . . — Big-work-item feature work-stream lifecycle (own copy-on-write project copy + own branch/tags + per-feature builds, no-merge-until-approved, trunk-merged-into-every-stream) (User mandate, — —)

Every BIG work item — a new feature OR a drastic/large fix — MUST be developed as its own isolated **feature work-stream**: its OWN full project copy in a sibling directory, its OWN feature branch, its OWN per-feature flashable builds + feature release tags, kept SEPARATE from the trunk until manually approved after testing, with the trunk regularly merged INTO every feature stream. This generalises Git-flow's feature/release/tag discipline (long-lived `main` trunk, one `feature/<slug>` per big item, no merge to trunk until reviewed-complete) and binds it to the project's parallel-stream (§11.4.58/§11.4.103), single-builder, and finite-test-device realities. The mandate (ALL hold):

(A) One big item = one feature work-stream in a sibling project copy. Each big work item gets its own complete project copy in a sibling directory (`<project>` , `<project>_<slug-1>` , `<project>_<slug-2>` , ...), the `<slug>` a lowercase identifier per §11.4.29. The trunk copy is the canonical `main` . Small changes stay on the trunk; only BIG items earn a stream (the threshold is the project's own per §11.4.35).

(B) Disk-feasible duplication is MANDATORY, naive copies FORBIDDEN. A feature-stream copy MUST be created by a near-zero-disk mechanism — on a copy-on-write filesystem (btrfs/XFS/ZFS/APFS) a reflink/CoW clone (`cp -a --reflink=auto` / filesystem snapshot) so only DIVERGED bytes consume disk; elsewhere a shared-object-store mechanism (e.g. `git worktree` + externalized build output + sparse where the build tolerates it). A full deep copy or full `git clone` per stream of a large working tree is FORBIDDEN where it would exhaust free space (§9 data safety + the project's host-disk reality per §11.4.35). The chosen mechanism MUST be verified to consume ~0 disk for a fresh stream (captured evidence per §11.4.69) before being relied on (§11.4.6 — no guessing about disk). Build output (`out/` -class) MUST be per-stream and excluded from the CoW clone; one shared compiler cache (ccache-class) across streams maximises queued- build hit-rate.

(C) Own branch + own tags, separate until approved. Each stream develops on its own `feature/<slug>` branch; per-feature flashable builds are marked with feature tags under a scheme that embeds the slug so the build is greppable across every repository (composes §11.4.151), e.g. `<base-version>-feat-<slug>[ -<iter>]` , the `<iter>` monotonic per §11.4.54. **The feature branch is NEVER merged to the trunk until the operator manually approves after testing** — approval requires the §11.4.40 full-suite retest GREEN on a clean target, then the §11.4.41 merge-first integration, then a §11.4.113 ff-only (no-force) merge to trunk.

(D) Trunk merged INTO every feature stream, regularly. The trunk MUST be merged INTO every live feature stream FREQUENTLY (at minimum after every trunk tag / daily) to prevent long-lived-branch conflict storms — `git merge` (never rebase a shared/tagged branch). A stream that has not been synced within the cadence is flagged stale.

(E) Feature-level branching AND tagging cascades to ALL submodules the moment a submodule is modified. The instant a stream modifies ANY submodule, that submodule gets its own `feature/<slug>` branch; per-feature tags on the main repo are mirrored on EVERY touched submodule at the same HEAD and ff-only-pushed to all owned remotes (§11.4.113 + the project's tag-cascade rule). Untouched submodules stay trunk-pinned (no branch/tag churn). Merge-to-trunk applies the same cascade. Where the duplication mechanism is `git worktree` , the experimental/disk- duplicating worktree-submodule path (each

worktree keeps a unique submodule copy; removal must be forced) is a documented hazard and the CoW-clone mechanism (each stream owns CoW-shared submodule git dirs) is preferred where available.

(F) Single-builder build queue + finite-device test queue. Where the host can run only ONE build at a time, feature builds MUST QUEUE (FIFO + advisory lock, composes §11.4.58 L1 / §12.7 / §12.8), never contend. Where live testing has a finite device count, per-device exclusive ownership (§11.4.119) serialises tests; streams beyond the device count queue. The build queue MUST garbage-collect idle streams' regenerable build output (§11.4.77) to keep disk within budget.

(G) Every change crosses the full quality gauntlet. Every change in any stream passes the extensive independent code review (§11.4.142/§11.4.125/§11.4.134)

- multi-angle impact research (§11.4.145) + anti-bluff captured-evidence testing (§11.4 / §11.4.5 / §11.4.69 / §11.4.107) — a feature stream is NOT a quality carve-out.

(H) Resumable, data-safe stream lifecycle. A stream registry (single source of "is this stream owed?", composes §11.4.116 / §11.4.147) tracks each stream's branch, base tag, build/out state, last-trunk-merge, and merge-approval state. Stream creation is atomic-on-success (clone to a `.partial` name, rename on success) and resumable (§11.4.84 quiescence check on resume). Retiring a stream requires a §9.2 pre-op backup + a guard that the feature work is merged/approved (unmerged work blocks auto-deletion). Host-session safety (§12) and the ≤60%-memory ceiling (§12.6) bound clone + build.

Honest boundary (§11.4.6): this lifecycle guarantees ISOLATION, disk-feasibility, greppable identity, and a controlled trunk ↔ feature merge discipline — it does NOT by itself prove any stream's work is correct (each change still crosses §11.4.108 runtime-signature + §11.4.40 retest). Classification: universal (§11.4.17) — the consuming project supplies its filesystem/CoW mechanism, sibling- dir + tag naming, build-queue + device-queue counts, and big-item threshold per §11.4.35.

Propagation gate `CM-COVENANT-114-167-PROPAGATION` (literal `11.4.167` across the consumer fleet) + recommended gates `CM-FEATURE-WORKSTREAM-COW-CLONE` (a feature stream is created by the verified near-zero-disk mechanism, NOT a full deep copy/ clone; per-stream out/ + shared cache; captured disk-cost evidence) + `CM-FEATURE-WORKSTREAM-NO-MERGE-UNTIL-APPROVED` (no feature → trunk merge without the approval marker + §11.4.40 GREEN) + `CM-FEATURE-WORKSTREAM-`

TRUNK-SYNC-CADENCE (every live stream's last-trunk-merge within the declared cadence) + CM-FEATURE-WORKSTREAM-SUBMODULE-CASCADE (every touched submodule carries the feature/<slug> branch + mirrored feature tag). Paired §1.1 meta-test mutations: strip the literal → propagation gate FAILs; create a stream via full deep copy → CM-FEATURE-WORKSTREAM-COW-CLONE FAILs; merge a feature to trunk with no approval marker → CM-FEATURE-WORKSTREAM-NO-MERGE-UNTIL-APPROVED FAILs; let a stream go stale past the cadence → CM-FEATURE-WORKSTREAM-TRUNK-SYNC-CADENCE FAILs; modify a submodule in a stream without its feature/<slug> branch/tag → CM-FEATURE-WORKSTREAM-SUBMODULE-CASCADE FAILs (gate-code = separate work item).

Composes §9 / §9.2 / §11.4.6 / §11.4.17 / §11.4.28 / §11.4.29 / §11.4.35 / §11.4.40 / §11.4.41 / §11.4.54 / §11.4.58 / §11.4.66 / §11.4.69 / §11.4.71 / §11.4.77 / §11.4.84 / §11.4.103 / §11.4.108 / §11.4.113 / §11.4.116 / §11.4.119 / §11.4.125 / §11.4.142 / §11.4.145 / §11.4.147 / §11.4.151 / §12 / §12.6 / §12.7 / §12.8 / §12.9 / §1.1.

Canonical authority: constitution submodule Constitution.md §11.4.167. Non-compliance is a release blocker. No escape hatch — no --full-clone-per-feature, --deep-copy-stream, --skip-cow-clone, --merge-feature-without-approval, --no-trunk-sync, --skip-submodule-cascade, --unbounded-parallel-builds, --shared-device-no-lock, --retire-unmerged-stream, --feature-stream-quality-carveout flag.

2026 06 23

§ . . — Exported-document
independent content + textual + full-visual
validation mandate (User mandate,
— —)

Forensic anchor (FACT, 2026-06-23). A "green" Mermaid export fix shipped RAW gantt-diagram source — the literal lines gantt, title, dateFormat, section, :done, — as plain readable text into user-facing PDFs, unreadable to any reader who expected a rendered diagram. The fix's validation grepped the HTML for class="mermaid" but NEVER opened the EXPORTED PDF to check whether raw diagram source had leaked into the body. The operator caught it (§11.4.138 operator-escape — the green suite missed it). Root lesson: generated/exported documents were never INDEPENDENTLY validated for the VISUAL + TEXTUAL correctness of the EXPORTED artifact itself — a grep of the source-or-intermediate HTML is NOT a check of what the reader actually receives. This is the

§11.4.108 SOURCE → ARTIFACT gap (the change present in source / intermediate, but the rendered bytes a reader opens are broken) materialised at the documentation-export layer.

The mandate. Every generated/exported document — HTML, PDF, DOCX, or any format in the §11.4.65 universal-Markdown-export scope (and the §11.4.153 four-format feature-Status set that adds DOCX) — MUST pass INDEPENDENT validation by a review agent structurally separate from the generator (§11.4.70/§11.4.20 dedicated subagent or a distinct human; NEVER the author/generator self-checking its own output — §11.4.92 self-evaluation PRECEDES and never satisfies; this is the §11.4.142 every-change-reviewed discipline applied to exported artifacts). The validation MUST be run on BOTH the source Markdown AND every exported artifact derived from it, ALWAYS — every export, not a sampled subset — across THREE layers, ALL of which MUST hold:

(1) CONTENT layer. The export faithfully carries the source's intent + data: nothing dropped, truncated, re-ordered into nonsense, or garbled. Every section, table, list, code block, figure caption, and link the source declares is present and complete in the export (a table truncated to its first page, a list collapsed, a section silently omitted, a link stripped of its target are all CONTENT findings). Cross-check the exported artifact against the source — present-and-complete, not merely "the file exists and is non-empty" (the §11.4.5/§11.4.69 captured-evidence quality bar, not a presence-only bar).

(2) TEXTUAL layer. The export is human-readable: NO raw markup, NO diagram source, NO unrendered code-fence body leaking as ordinary body text. Specifically forbidden as body text in the exported artifact: raw Mermaid diagram source (`mermaid`, `gantt`, `graph`, `flowchart`, `sequenceDiagram`, `classDiagram`, `stateDiagram`, `erDiagram`, `pie`, `gitGraph`, `journey`, `dateFormat`, `section ...`, `:done`, `/ :active`, `/ :crit`, task lines, and the closing-fence residue), raw HTML tags rendered as visible text, raw LaTeX/math source where math rendering was intended, and any triple-backtick code-fence whose intended rendering (diagram / formatted block) was supposed to replace the source. The 2026-06-23 raw-gantt-in-PDF failure is the canonical TEXTUAL finding. Detection reads the exported artifact as text (e.g. `pdftotext` on the PDF, the rendered DOM/text of the HTML, the extracted document text of the DOCX) and asserts the forbidden tokens do NOT appear as body text.

(3) FULL VISUAL layer. Embedded diagrams render as IMAGES, not as source; layout is intact; no overlapping / cut-off / clipped / garbled / collapsed / blank content. Verified by RENDERING the export and inspecting the rendered result — never by inspecting the source or an intermediate alone: e.g. `pdftotext` to catch raw diagram source leaking as text (composes layer 2), `pdfimages` (or the format's image-enumeration) to CONFIRM the intended diagrams are present as rendered raster/vector images (a diagram that should be an image but produces zero embedded images is a VISUAL finding — the source leaked instead of rendering), and `pdftoppm` → OCR (§11.4.117/§11.4.107(12), confidence floor + ROI) to CONFIRM the rasterised page carries human-readable visual content and to re-catch any raw-source-as-pixels. Captured evidence per §11.4.5/§11.4.69/ §11.4.107 (`pdftotext_body.txt` , `pdfimages_manifest.txt` , `ocr_pages/` , the per-finding pinpoint — which page/region/token, expected vs actual) accompanies every verdict.

(4) Anti-bluff — the validator is itself anti-bluff. A rubber-stamp "the export looks fine" is NOT validation (§11.4 / §11.4.1 — a rubber-stamp validator is a PASS-bluff at the export-validation layer). The analyzer/validator is self-validated with a golden-good / golden-bad fixture pair per §11.4.107(10): a golden-good export (clean rendered diagram + faithful content) MUST PASS, and a golden-bad export (a PDF/HTML/DOCX deliberately seeded with raw Mermaid `gantt` source as body text, plus a truncated table) MUST FAIL — a validator that passes its golden-bad fixture is a bluff gate and is investigated per §11.4.102 before any further use. The validation iterates to a clean GO per §11.4.134: any finding in ANY of the three layers (BLOCKING / nit / warning) re-arms the loop — the generator/author fixes it (e.g. pre-renders the diagram to an image before export), the export is regenerated, and the validator RE-VALIDATES, until GO with ZERO findings and ZERO warnings, every verdict backed by captured rendered-artifact evidence. A generated document that ships WITHOUT this independent content+textual+visual validation, or that contains raw diagram source / unreadable garble / dropped content in the exported artifact, is a §11.4 PASS-bluff at the documentation layer.

Honest boundary (§11.4.6). This validation confirms the EXPORTED artifact a reader opens carries the source's content faithfully, is textually readable, and renders its visuals as images — it does NOT prove the source's content is itself correct (that rests on the source's own §11.4.142 review + §11.4.135 regression guards) and does NOT replace the §11.4.65 mtime-parity export-sync gate (which proves the exports are FRESH, not that they are READABLE — §11.4.168 is the

READABILITY/FIDELITY layer the mtime gate cannot see). A FLAG_SECURE / DRM-blanked / sink-blanked rendering surface that cannot be rasterised documents the gap per §11.4.112 and uses the §11.4.117 proxy oracle — never a faked visual PASS. "The HTML had `class="mermaid"` " is a source-side check, NOT an exported-artifact check (the exact 2026-06-23 omission); " `pdftotext` shows no `gantt` source AND `pdfimages` shows the expected diagram image AND OCR reads human text" is the exported-artifact check.

Classification: universal (§11.4.17) — a platform-neutral exported-document-fidelity discipline reusable by ANY project that generates/exports documents; the consuming project supplies its concrete pre-render pipeline, exporters, and rendered-artifact validators (text-extraction / image-enumeration / rasterise → OCR) per §11.4.35. Composes §11.4.65 (the export-scope + freshness layer this READABILITY layer sits atop) / §11.4.73 (project-wide MD/HTML/PDF/DOCX styling — styled exports MUST still pass the three layers) / §11.4.107 (the self-validated-analyzer + liveness/OCR validation techniques) / §11.4.117 (CV/OCR pixel-oracle for rasterised-page validation) / §11.4.135 (a fixed export-fidelity defect registers a permanent regression guard) / §11.4.138 (the 2026-06-23 operator-escape that motivated this anchor) / §11.4.142 (independent every-change review applied to exports) / §11.4.159 (window-scoped recording + vision-validation discipline, the recording analogue) / §11.4.163 (the media-validation pipeline — §11.4.168 is its document-export sibling) / §11.4.165 (the independent verification agent — §11.4.168(c) DOCS layer is its export-fidelity refinement) / §1.1.

Propagation gate `CM-COVENANT-114-168-PROPAGATION` (literal `11.4.168` across the consumer fleet) + recommended gate

`CM-EXPORTED-DOC-VISUALLY-VALIDATED` (every exported document in the §11.4.65 scope carries an independent-validation marker — produced by a validator structurally separate from the generator — asserting the three layers PASS on the exported artifact, with captured rendered-artifact evidence; the validator is self-validated by a golden-good/golden-bad fixture pair) + paired §1.1 meta-test mutation (strip the literal `11.4.168` → propagation gate FAILs; produce a PDF/HTML/DOCX containing raw diagram source — e.g. raw Mermaid `gantt` lines — as body text and assert the gate catches it → `CM-EXPORTED-DOC-VISUALLY-VALIDATED` FAILs; feed the validator its golden-bad fixture → self-validation FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.168. Non-compliance is a release blocker regardless of context. No escape hatch — no `--skip-visual-validation`, `--raw-source-in-pdf-ok`, `--textual-check-suffices`, `--self-validate-export`, `--diagram-source-as-text-ok` flag exists.

2026 06 25

§ . . — **Mandatory comprehensive test-type coverage with anti-bluff captured evidence (User mandate, — —)**

Forensic anchor — verbatim user mandate (2026-06-25):

"MANDATORY test types EACH project MUST have with coverage aiming to 100% without false results, rock solid evidence produced covered with hard physical proofs and zero bluff: unit, integration, e2e, full automation, challenges (challenges submodule), helix qa (helix qa submodule with proper tests banks (suites) and comprehensive fully autonomous qa sessions), ddos tests, security tests, stress and chaos tests, concurrency (atomicity) tests, race conditions or deadlocks tests, memory tests, benchmarking tests."

Every project under this Constitution MUST be covered by the following CLOSED, ENUMERATED set of test types, each driven to **as close to 100% coverage as the domain permits**, with every PASS backed by **rock-solid captured PHYSICAL evidence** (§11.4.5 / §11.4.69 / §11.4.107) and **zero false results, zero bluff of any kind** (§11.4 / §11.4.1 / §11.4.6). A green result with no falsifiable captured evidence is a §11.4 PASS-bluff regardless of which layer reports it. §11.4.169 is the strict, explicitly-enumerated expansion of §11.4.27 (no-fakes-beyond-unit + 100%-test-type coverage) — it freezes the canonical test-type SET and binds each type to captured evidence and, where named, to the owned submodule that supplies it.

The mandatory test-type set (each REQUIRED where the domain warrants it; an honest §11.4.3 SKIP-with-reason — topology / hardware / credential genuinely absent — is the ONLY permitted absence, NEVER a silent gap):

1. **Unit** — isolated logic; the ONLY layer where mocks / stubs / fakes / placeholders are permitted (§11.4.27(A)).

2. **Integration** — real, fully-wired components against real infrastructure (no fakes beyond unit), infra booted on-demand via the **containers** submodule (§11.4.76 / §11.4.161 rootless).
3. **End-to-end (e2e)** — the full user journey across the real System, captured-evidence per §11.4.5 / §11.4.107.
4. **Full-automation** — fully autonomous, re-runnable end-to-end with NO manual intervention after start (§11.4.25 / §11.4.52 / §11.4.98), deterministic across N iterations (§11.4.50).
5. **Challenges** — the **challenges** submodule (`basic-digital/challenges`): anti-bluff Challenge banks that score PASS only on positive captured evidence (§11.4.27(B) / §11.4.4(b) layer 4).
6. **HelixQA** — the **helix_qa** submodule with proper written **test banks (suites)** for every application / service / platform AND comprehensive **fully-autonomous QA sessions** driving every registered bank with captured wire evidence per check (§11.4.27).
7. **DDoS / load-flood** — sustained adversarial traffic; the System refuses cleanly OR degrades gracefully, never collapses; captured throughput / latency / error-rate evidence.
8. **Security** — authn / authz, injection / taint, secret-leak (§11.4.10), transport + crypto, dependency CVEs, default-deny enforcement; composes the **security** submodule.
9. **Stress + Chaos** — sustained load + concurrent contention + boundary inputs + failure-injection (process-death / network-fault / input-corruption / resource-exhaustion / state-corruption) with categorised recovery + captured evidence (§11.4.85).
10. **Concurrency / atomicity** — correctness under concurrent callers; no lost updates; transactional atomicity proven; idempotency under replay.
11. **Race-condition / deadlock** — race detector + lock-order / deadlock analysis under contention; no blocking operation inside a shared-lock region (the §1 development principles); captured detector output is the evidence.
12. **Memory** — leak census over soak, peak-RSS ceilings (e.g. the iOS-NEPacketTunnelProvider-class budget), allocation / fragmentation profile; no unbounded growth across a 24h / N-iteration soak.

13. **Benchmarking / performance** — p50 / p95 / p99 latency + throughput + resource cost vs a recorded baseline; a regression vs baseline is a finding, not noise.

Four-layer enforcement per §11.4.4(b) of the suite itself: pre-build gate (each required type present + executable + parseable for the domain), post-build, on-device / runtime where applicable, and a paired §1.1 meta-test mutation per type (strip a type's evidence-capture or its anti-bluff assertion → the gate FAILs). The per-project coverage ledger (§11.4.25 / §11.4.52) **MUST** classify every feature × test-type × evidence-state; a missing required type, or an

`OPERATOR_ATTENDED_ONLY` row, is a release blocker until promoted or honestly SKIP-justified (§11.4.3 / §11.4.52).

Composes / strengthens §11.4.4 / §11.4.5 / §11.4.6 / §11.4.25 / §11.4.27 (its strict enumerated expansion) / §11.4.50 / §11.4.52 / §11.4.69 / §11.4.76 / §11.4.85 / §11.4.98 / §11.4.107 / §11.4.123 / §11.4.135 / §11.4.146 / §1.1. Classification: universal (§11.4.17) — the consuming project supplies its concrete per-type harnesses, infra, and calibrated thresholds per §11.4.35.

Pre-build gate `CM-COVENANT-114-169-PROPAGATION` (literal `11.4.169` across the consumer fleet) + recommended gate `CM-MANDATORY-TEST-TYPES-COVERED` (every required test type present-or-honestly-SKIPPED per feature, each PASS citing captured evidence) + paired §1.1 meta-test mutation (strip a required type or drop its evidence citation → gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.169. Non-compliance is a release blocker regardless of context. No escape hatch — no `--skip-test-type`, `--partial-coverage-ok`, `--evidence-optional`, `--challenges-not-applicable`, `--skip-helixqa`, `--no-chaos`, `--no-ddos`, `--skip-memory-test`, `--skip-race-detector`, `--bench-optional` flag exists.

[COLLISION-REDIRECT 2026-08-18] The 2026-06-25 HelixTranslate canonical translation pipeline mandate previously carried §11.4.140 here has been re-minted to **§11.4.255** to resolve the pre-existing double-mandate collision surfaced by §11.4.227(B) block-integrity audit. The other §11.4.140 mandate (Universal action-prefix system, 2026-06-09) at line ~9114 is unchanged and remains the sole occupant of the §11.4.140 anchor number. Operator-approved re-mint via AskUserQuestion on 2026-08-15.

[COLLISION-REDIRECT 2026-08-18] The 2026-06-25 Mandatory independent per-language translation review mandate previously carried §11.4.141 here has been re-minted to **§11.4.256** to resolve the pre-existing double-mandate collision surfaced by §11.4.227(B) block-integrity audit. The other §11.4.141 mandate (Token-efficiency, 2026-06-09) at line ~9233 is unchanged and remains the sole occupant of the §11.4.141 anchor number. Operator-approved re-mint via AskUserQuestion on 2026-08-15.

§11.4.170 — Device-independent host-side rendered-UI visual-proof mandate (User mandate, 2026-06-25). Verbatim operator escalation set (2026-06-25): "what this should represent, this ugliness?!?!? UI we see does not purpos or sense at all! it is messed up with an giant button!" + "Eye candy stunning funcitonal UI and UX MUST BE prooven with full automation tests which 'see' everything visually with no issues, false results or bluff of any kind!!! So far, anti-bluff policy seems completely ignored!" Every change to any user-facing UI surface (screen / component / view / widget / layout / theme on any platform — Android-Compose, iOS-SwiftUI/UIKit, web, desktop, TUI) MUST be proven by **device-independent host-side RENDERED PIXELS** before it may be claimed correct: the actual component rendered to an image ON THE HOST (NO device, NO emulator, NO running app required) via a host-render harness (Compose → Roborazzi/Paparazzi on the JVM; web → Playwright/Storybook snapshot; SwiftUI → snapshot-testing; equivalent per stack), for EVERY relevant `screen × state × {light,dark}`

theme combination, with the produced PNG validated by BOTH (i) **image-diff against an approved golden** (per-pixel or perceptual-diff PASS/FAIL) AND (ii) an **OCR / vision oracle** that reads the rendered text + labels + control bounds to assert legibility + NO element overlap / NO label-over-label / NO clipping / NO off-screen control / NO collapsed-or-giant unbounded widget (the §11.4.162 layout invariants, here PROVEN on real pixels not asserted on tokens). **The forbidden bluff (forensic anchor, FACT 2026-06-25):** UI work was "validated" by VALUE-EQUALITY unit tests (asserting a hex colour string, an `sp / dp` size integer, a design-token field equals an expected constant) that NEVER RENDERED A SINGLE PIXEL — every such test was GREEN while the operator opened the running screen and saw a broken, unbounded giant-button layout. A value-equality / token-equality / property-assertion UI test is therefore FORBIDDEN as the *proof* a UI is correct (it may SUPPLEMENT — confirming a token's wiring — but NEVER SUBSTITUTE for the rendered-pixel proof); a "green" UI suite that contains no rendered-pixel artefact for the changed surface is a §11.4 PASS-bluff at the visual layer, severity-equivalent to a metadata-only PASS. The mandate (ALL hold): (1) **rendered-pixel proof per screen × state × theme**, device-independent, runnable in CI on EVERY commit with zero hardware (that is the whole point of host-side rendering — it is NOT subject to device availability); (2) **dual oracle** — golden image-diff AND OCR/vision content+layout oracle, both required, each PASS citing the captured PNG + diff artefact path per §11.4.5/§11.4.69; (3) **value-only UI claims FORBIDDEN as proof** — a UI change with no rendered-pixel artefact for its surface is a finding (§11.4.4 test-interrupt → render → fix → re-render GREEN); (4) **self-validated analyzer** — the image-diff + OCR oracle is golden-good/golden-bad calibrated per §11.4.107(10) (a golden-bad fixture seeded with overlapping labels / a clipped control / an unbounded button MUST FAIL, else the analyzer is itself a bluff gate), thresholds calibrated on the project's OWN fixtures not hardcoded from literature (§11.4.6); (5) **"device offline" is NEVER a valid skip** for this proof — host-render is precisely the device-independent path; an honest §11.4.3 SKIP applies ONLY where the platform genuinely has no host-render harness (a tracked migration item, never a silent gap or a value-only fallback). Honest boundary (§11.4.6): host-render proves the COMPONENT renders correctly device-independently on every commit — it does NOT prove the running app on real hardware (that remains §11.4.153/§11.4.158/§11.4.159/§11.4.160's live-device recording + read-the-screen layer, which §11.4.170 COMPLEMENTS not replaces — the two layers are both mandatory: host-render for per-commit device-free coverage, live-recording for end-to-end on-device confirmation), does NOT replace

§11.4.162's OpenDesign token/theme governance (§11.4.170 PROVES the rendered RESULT of those tokens), and is a DISTINCT artifact class from §11.4.168's exported-document visual validation. Classification: universal (§11.4.17) — the consuming project supplies its concrete host-render harness, golden corpus, screen×state×theme matrix, and OCR/vision thresholds per §11.4.35. Composes §11.4.3/§11.4.5/§11.4.27/§11.4.40/§11.4.69/§11.4.107/§11.4.108/§11.4.117/§11.4.153/§11.4.158/§11.4.159/§11.4.160/§11.4.162/§11.4.163/§11.4.168/§11.4.169/§1.1. Propagation gate CM-COVENANT-114-170-PROPAGATION (literal 11.4.170) + recommended gate CM-HOST-RENDERED-UI-VISUAL-PROOF (every changed UI surface carries a device-independent host-rendered PNG per screen×state×theme, dual-oracle-validated, run on every commit; a UI change proven only by value/token-equality assertions FAILs) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; accept a UI change whose only proof is a value-equality assertion / a golden-bad overlapping-label fixture that PASSes → CM-HOST-RENDERED-UI-VISUAL-PROOF FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule Constitution.md §11.4.170. Non-compliance is a release blocker. No escape hatch — no --value-assertion-proves-ui, --skip-rendered-pixel-proof, --token-equality-suffices, --device-offline-skip-visual, --single-theme-only, --unvalidated-image-diff-OK, --no-ocr-layout-oracle flag.

§ . Host-session safety – directly OR indirectly signing the user out is FORBIDDEN

¹² ¹ Every script, test, helper, and AI agent governed by this Constitution MUST respect host-session safety. Non-compliance is a release blocker.

§ . Forbidden operations – directly OR indirectly

1. **Suspending the host:** systemctl suspend, pm-suspend, loginctl suspend, DBus org.freedesktop.login1.Suspend, GNOME idle-suspend, lid-close handler.

2. **Hibernating / hybrid-sleeping:** any `Hibernate` / `HybridSleep` / `SuspendThenHibernate` method.
3. **Logging out the user:** `loginctl terminate-session` , `pkill -u <user>` , `systemctl --user --kill` , anything that signals `user@<uid>.service` .
4. **Unbounded-memory operations** inside `user@<uid>.service` cgroup. Any single command expected to exceed ~4 GiB RSS MUST be wrapped in a bounded execution scope (e.g. `bounded_run` from a project helper library that wraps `systemd-run --user --scope -p MemoryMax=...`).
5. **Programmatic rkill toggles, lid-switch handlers, or power-button handlers** — these cascade into idle-actions.
- 6.^{1 2} **Disabling systemd-logind, GDM, or session managers** "to make things faster" — even temporary stops leave the system unable to recover.

§ . Required safeguards

Every script in this project that performs heavy work (build, transcription, model inference, large compression, multi-GB git operations) MUST:

1. Source the project's host-safety helper library at the top.
2. Call its pre-flight check and **abort if it fails**.
3. Wrap any subprocess expected to exceed ~4 GiB RSS in a bounded execution scope so the kernel OOM killer is contained to that scope and cannot escalate to user.slice.
- 4.^{1 2 3} Cap parallelism (`-j`) to fit available RAM (estimate per-job-peak-RSS and divide into the budget).

§ . Container hygiene

Containers (Docker / Podman) the project owns or relies on MUST:

1. Declare an explicit memory limit (`mem_limit` / `--memory` / `MemoryMax`).
2. Set `OOMPolicy=stop` in their systemd unit to avoid retry loops.
3. Use exponential-backoff restart policies, never immediate retry.
4. Be clean-slate destroyed and rebuilt after any host crash or session loss so stale lock files don't keep producing failures.

§ . Memory-Budget Ceiling – % MAXIMUM

Forensic anchor — direct user mandate:

"First make sure that whatever we do through our procedures related to this project MUST NOT use more than 60% of total system memory! All processes MUST be able to function normally!"

The mandate. Project procedures MUST NOT use more than **60% of total system RAM**. The remaining 40% is reserved for the operator's other workloads so the host can keep serving them while project work proceeds.

The protections:

1. `HOST_SAFETY_MAX_MEM_PCT` defaults to 60.
2. `HOST_SAFETY_BUDGET_GB` is computed at source-time from `MemTotal × MAX_PCT/100`, in GiB.
3. Bounded execution scopes clamp `MemoryMax` down to the budget if the caller asks for more.
4. The build script's parallelism is computed as `min(nproc, floor(budget_gb / per_job_peak_rss_gb))`.
5. Heavy work MUST be wrapped in a bounded execution scope so the kernel OOM-kills only the scope — `user@<uid>.service` stays alive.

No escape hatch. §12.6 has NO operator-facing override flag. The cap exists for the operator's own protection; bypassing it is the bluff the §11.4 covenant specifically prohibits.

§ . Continuation document – sacred invariant

Forensic anchor — direct user mandate:

"during any work we perform, during Phases implementation, debugging and fixing, during ANY effort we have the Continuation document MUST BE maintained and it MUST NOT BE out of sync with current work we are doing! If for any reason we stop our work, we MUST BE able to continue any time, with current work, exactly where we have left off and from any CLI agent or any LLM model we chose! Nothing can be broken or faulty in maintained Continuation document!"

The mandate. A single, canonical, machine-readable handoff document — `docs/CONTINUATION.md` — must always reflect the live state of the project. Any agent (human, Claude Code, Cursor, Aider, Codex, Gemini CLI, any future LLM) must be able to resume work **exactly where the previous session left off** by reading this single file.

Mandatory protections:

1. `docs/CONTINUATION.md` **MUST exist** at the project root. Its absence is a release blocker.
2. **Every non-trivial state change** MUST update this document in the same commit as the work itself.
3. **Top-of-file timestamp** is updated on every edit. Stale timestamps trigger gate failure.
4. **Section §3 "Active work"** lists every IN PROGRESS / BLOCKED item with enough detail that any agent can resume without conversation context — concrete commands, file paths, monitor IDs.
5. **Section §0 "How to use this document"** contains the verbatim resumption prompt — a single block any operator can paste into any CLI agent.
6. **Document MUST be self-contained.** No hyperlinks to ephemeral external systems as the only source of truth.

No escape hatch. §12.10 has NO operator-facing override flag for the existence requirement. The discipline exists for the operator's own protection.

§ . — Maximal dynamic resource utilization for containerized builds (User mandate, — —)

Forensic anchor — verbatim user mandate (2026-07-03):

"Use full workstation system capacity — all resources to build and test! Give maximal number of cores and memory! Use dynamic capabilities! No limiting to -j2 anymore! ... MUST NOT break the system, create bottlenecks, or get stuck! ... MUST BE now part of constitution Submodule! On each System we use, aim for maximal use of available resources for maximal efficiency!"

The reconciliation (the load-bearing logic). The existing memory-and-parallelism caps were born from a specific, forensically-captured failure mode: **heavy build work running inside the user login session's cgroup (`user@<uid>.service` / `user.slice`) overran memory and the kernel OOM-killer escalated to the session itself, SIGKILLing the operator's desktop, terminals, and the agent** (the §12.6 forensic anchor: "three consecutive session-loss SIGKILLS" during a bare-metal `m -j5` ; consuming-project instantiations extend this — a bare-metal AOSP `m -j` was hard-capped at -j2 because `soong_build` alone holds a large constant RSS inside `user.slice`). Those caps — **§12.6's 60%-of-total-RAM ceiling** and the consuming project's bare-metal `-j` cap — are CORRECT and REMAIN in force for **any heavy work that executes inside `user@<uid>.service` / `user.slice`** , because such work shares the session's OOM fate.

A build that runs in **its OWN container cgroup with its OWN `MemoryMax`** (the rootless-Podman containerized build path — generically "the sanctioned containerized build") is a **fundamentally different accounting domain**: it is NOT a child of `user.slice` , so the kernel OOM-killer and any userspace OOM watchdog (systemd-oomd's `user.slice` PSI watch) act on the **container's own cgroup FIRST** — the container OOM-kills ITSELF, it does not escalate to and cannot SIGKILL the operator session. Therefore a high memory / CPU / `-j` envelope in the containerized path is **host-SAFE by construction** — the exact hazard the §12.6 ceiling guards against does not reach the session across the cgroup boundary.

The mandate. For the **containerized build path only**, the project **MUST** use the **maximal SAFE fraction** of the host's real capacity, computed **dynamically per host** — NOT a fixed low `-j` count, NOT clamped to the §12.6 60% figure (which governs user.slice-resident work, a different domain). Concretely (ALL hold): (1) **Auto-detect host capacity, never hardcode** — read this host's real `nproc`, `MemTotal`, `SwapTotal`, and the hard process/thread rlimit (`RLIMIT_NPROC` / `ulimit -Hu`) at launch; every number driving the envelope **MUST** be measured on the running host, never assumed (§11.4.6 no-guessing); a launcher that prints a build envelope it did not measure is a §11.4.6 violation. (2) **Reserve explicit host headroom (cores AND RAM)** — the maximal-safe fraction is "all capacity MINUS a reserved slice kept free for the operator's live session + other workloads"; the headroom **MUST** be explicit and captured in the printed envelope; this preserves the SPIRIT of §12.6 (the operator's session keeps functioning) via a reserved-headroom mechanism, WITHOUT importing §12.6's 60% number — the container's cgroup, not a user.slice budget line, bounds it. (3) **`-j` scales nproc-awarely, not just memory-awarely** — the parallel-job count **MUST** respect BOTH the memory governor (`floor((container_mem_gb - build_constant_overhead_gb) / per_job_peak_rss_gb)`) AND the process/thread rlimit, because each `-j` unit spawns compiler/linker worker THREADS that count against `RLIMIT_NPROC` and exceeding it causes `EAGAIN` "fork: retry" storms that stall/wedge the build (the "never get stuck" clause); where the login's hard `nproc` is too low for the full memory-and-CPU-derived `-j`, the launcher **MUST** either (a) run at the proven-safe lower `-j` under the rlimit, OR (b) on a privileged path raise the rlimit (e.g. `prlimit`) and drop back to the unprivileged uid (e.g. `setpriv`) before launching so the container keeps its rootless identity; the launcher **MUST** print which mode it chose (§11.4.6 — the envelope is captured evidence). (4) **The caps REMAIN for the bare-metal / user.slice path** — §12.6's 60% ceiling and the consuming project's bare-metal `-j` cap are NOT relaxed for any work inside `user@<uid>.service`; §12.11 grants the maximal envelope **ONLY** across the container-cgroup boundary; a bare-metal heavy build is still forbidden from claiming the §12.11 envelope (that would reproduce the original session-loss incident); §12.11 does NOT weaken §12.6, it SCOPES it. (5) **Never break, never bottleneck, never get stuck** — the envelope **MUST** leave the reserved host headroom genuinely free (verifiable post-launch: session memory-available stays above the reserved floor), **MUST NOT** drive swap-thrash inside the container (bound `MemorySwapMax` so the container hits its own OOM before pathological

swap), and MUST NOT exceed the process rlimit (per clause 3); "maximal" means maximal-SAFE, not maximal-reckless — an envelope that wedges the host is a §12.11 violation, not compliance. (6) **Evidence-based tuning, no guessing (§11.4.6 + §11.4.24)** — the chosen envelope is driven by captured host metrics + the build's own resource profile (§11.4.24 build-resource stats); "use a bigger `-j`" without captured headroom evidence proving the host can absorb it is a §11.4.6 + §12.11(5) violation; when in doubt, the reserved-headroom floor wins.

Reference implementation (consuming project). A thin DYNAMIC wrapper over the sanctioned containerized build (ATMOSphere: `scripts/build_maxres.sh` over `scripts/build_containerized.sh`) reserves a fixed core count + a RAM floor (e.g. 72% of MemTotal to the container, clamped so a fixed-GiB slice stays free for the session), computes `JOBS_MEM = floor((container_mem_gb - constant_overhead_gb) / per_job_peak_rss_gb)`, and governs `-j` against the real hard `nproc : unprivileged` → the proven-no- EAGAIN job count; privileged (root re-exec raises `RLIMIT_NPROC` via `prlimit` then `setpriv` -drops back to the unprivileged uid) → full `-j = min(cpus, JOBS_MEM)`; `--dry-run` prints the full envelope and launches nothing. This is the §11.4.35 project instantiation; the mechanism above is universal.

Composes with §12.1 / §12.2 / §12.3 (host-session safety — §12.11 is a strengthening within, not a relaxation of, host safety) / **§12.6** (the 60% ceiling — §12.11 SCOPES it to user.slice-resident work, does NOT weaken it) / §12.10 (the launcher's envelope is a live-state anchor for the resumption file) / §11.4.6 (no-guessing — every envelope number measured) / §11.4.24 (build-resource stats feed the tuning) / §11.4.133 (target/hardware safety — the reserved-headroom + rlimit governors are the host-safety equivalent) / §9.2 (the containerized path is the non-destructive default). The consuming project supplies its containerized-build path, cgroup mechanism, and per-host detection per §11.4.35.

Classification: universal (§11.4.17) — the reconciliation logic (caps bind user.slice-resident work; the containerized cgroup is separately accounted so it may claim the maximal safe fraction) is platform-agnostic and true on every host. Propagation gate **CM-COVENANT-12-11-PROPAGATION** (literal anchor `12.11` present across the consumer fleet's CLAUDE.md / AGENTS.md / QWEN.md / GEMINI.md) + recommended gate **CM-MAXRES-DYNAMIC-BUILD** (the containerized-build launcher (a) auto-detects `nproc` / `MemTotal` / `RLIMIT_NPROC`, (b) reserves explicit host headroom, (c) governs `-j` against BOTH the memory

model AND the process rlimit, (d) prints the chosen envelope, (e) does NOT hardcode a fixed low `-j` for the containerized path AND does NOT relax §12.6 for any user.slice path) + paired §1.1 meta-test mutations (strip the `12.11` literal → propagation gate FAILs; replace the dynamic detection with a hardcoded `-j2` in the containerized launcher OR strip the rlimit-aware `-j` governor → `CM-MAXRES-DYNAMIC-BUILD` FAILs; gate-code = a separate tracked work item).

Canonical authority: constitution submodule `Constitution.md` §12.11. Non-compliance is a release blocker regardless of context. No escape hatch — no `--fixed-jobs1,2`, `--no-dynamic-detection`, `--ignore-host-capacity`, `--relax-user-slice-cap`, `--bare-metal-maximal` flag exists.

2026 07 07

§ . — Process/thread-limit
(`RLIMIT_NPROC`) awareness for parallel subagent/
multi-process work (User mandate,
— —)

Forensic anchor — verbatim user mandate (2026-07-07):

"Add these ulimit details and instructions to the constitution so they are mandatory for this and all projects!"

The limit. Heavy parallel subagent / multi-process work is bounded by the OS per-user process/thread limit `ulimit -u ( RLIMIT_NPROC )`, which on Linux counts ALL threads AND processes owned by the real UID — NOT a per-process limit. The ambient tooling fleet already running in a session (MCP servers, language runtimes — tokio/V8/libuv worker pools, `mongod` / `chrome` / `prisma` -class MCP servers, local LLM servers) consumes a large FIXED fraction of this budget before any project work even starts. **FORENSIC FACT (2026-07-07):** a host configured with `ulimit -u 4096` sat at ~3900–4005 threads from the ambient MCP fleet alone, leaving under 150 threads of headroom; dispatching a 4th Go-heavy subagent (the Go runtime sets `GOMAXPROCS` to `nproc`, spawning dozens of OS threads per invocation) hit `EAGAIN` / failed to create new OS thread (`errno=11`) and crashed tooling. Stopping three non-essential MCP servers (`mongod` / `chrome` / `prisma`) freed approximately 1345 threads (`3827` → `2482`), restoring headroom.

Mandate — treat thread-limit exhaustion as a §12 host-safety event. BEFORE scaling to N parallel subagents / processes, the agent MUST check thread headroom and treat thread-limit exhaustion as a §12 host-safety event that YIELDS UNCONDITIONALLY — the same standing as §12.6 (memory) and §12.7 (AOSP `-j cap`). This is an ORTHOGONAL exhaustion axis to §12.6's memory ceiling: a host can have abundant free RAM while its thread/process budget is fully exhausted, and vice versa — both MUST be checked independently before scaling parallelism.

Diagnosis (mandatory before scaling). `ulimit -u` (soft limit) + `ulimit -Hu` (hard limit) + the LIVE count `ps -L --no-headers -u "$USER" | wc -l` (threads, not just processes — `RLIMIT_NPROC` counts threads); headroom = soft limit minus live count. **Signals of exhaustion** (never guessed, always captured per §11.4.6): `EAGAIN`, `fork: retry: Resource temporarily unavailable`, `failed to create new OS thread`, or `cannot allocate memory returned from a fork / clone /thread-spawn call` — this is DISTINCT from OOM (memory can be abundant while the thread/process budget is exhausted; the diagnostic MUST NOT conflate the two). When headroom is low, the agent MUST SERIALIZE (run fewer concurrent subagents/processes) and avoid Go-heavy / `GOMAXPROCS=nproc` -class spikes — dispatching parallel work blind at near-limit headroom is forbidden; a subagent that hits `EAGAIN` mid-operation (e.g. mid- `git push`) risks incomplete/corrupted state (§9.2 data-safety risk).

Raising the limit — three techniques (cite the mechanism used, never guess which applies). (a) **No-restart, on a running process:**

`prlimit --pid <PID> --nproc=<soft>:<hard>` (requires root — e.g. via `su -c`) to raise a HARD limit above the process's current hard ceiling; a running process's new children inherit the raised limit going forward. (b) **Persistent, system-wide:** `/etc/security/limits.conf` (`<user> hard nproc <N> + a matching soft line`) or a systemd unit's `LimitNPROC=<N>` , taking effect on the NEXT login/session (not the current one). (c) **Self-raise, within the current session:** any process may raise its own SOFT limit up to its existing HARD limit via `ulimit -u <N>` where $N \leq \text{hard}$ — it can NEVER exceed the hard ceiling without root-assisted method (a). **Caveat (§11.4.6 — do not assume scope):** `su -c "ulimit -u N"` sets the limit ONLY inside that transient `su` subshell — it does NOT retroactively affect an already-running session; a running process keeps the `RLIMIT_NPROC` value captured at its own launch time.

Freeing threads — operator-authorized only (§11.4.122). Stopping non-essential MCP servers / background tooling frees the ambient fraction of the budget consumed before project work starts — but this **REQUIRES** explicit operator authorization per §11.4.122 (no-silent-removal); the agent **MUST NOT** autonomously kill the operator's tooling without asking first. **CAUTION — self-kill hazard (forensic FACT):** `pkill -f <pattern> / pgrep -f <pattern>` matches against the FULL COMMAND LINE, including the killer's own invocation if the search pattern string happens to appear in it — this can cause an unintended self-kill of the very process issuing the kill command. Any cmdline-pattern-based kill list **MUST** explicitly exclude the caller's own `$$` / `$PPID` (and the conductor's PID, when dispatched as a subagent) before executing.

Composes with §12 (host-session safety umbrella) / §12.6 (memory-budget ceiling — the ORTHOGONAL second exhaustion axis alongside this thread/process-count axis) / §12.7 (bare-metal AOSP `-j cap` — thread exhaustion is a DISTINCT failure mode from the memory-driven cap that section addresses) / §11.4.58 (parallel-development PWU pipeline — the concurrency this section bounds) / §11.4.103 (continuous parallel-stream working routine — the ≥ 3 -background-stream target is bounded by available thread headroom, not merely memory) / §11.4.101 (autonomous-decision-over-blocking — safe, reversible decisions, e.g. serializing dispatch, are made without operator involvement) / §11.4.122 (no-silent-removal-of-existing-components — freeing ambient tooling threads requires operator authorization) / §9.2 (absolute data safety — a subagent crashed mid-git-operation by `EAGAIN` is a data-safety risk this section exists to prevent).

Classification: universal (§11.4.17) — `RLIMIT_NPROC` / `ulimit -u` is a Linux/POSIX kernel-level per-UID accounting mechanism, platform-agnostic and true on every host running this Constitution's parallel-subagent/multi-process disciplines; the consuming project supplies its concrete ambient-tooling roster and `ulimit` values per §11.4.35. Propagation gate `CM-COVENANT-12-12-PROPAGATION` (literal anchor `12.12` present across the consumer fleet's `CLAUDE.md` / `AGENTS.md` / `QWEN.md` / `GEMINI.md`) + recommended gate `CM-NPROC-HEADROOM-CHECK` (parallel-dispatch orchestration checks `ulimit -u` + live thread count **BEFORE** scaling concurrency; a cmdline-pattern kill helper excludes the caller's own PID) + paired §1.1 meta-test mutation (strip the `12.12` literal → propagation gate FAILs;

strip the headroom-check call before a parallel-dispatch scale-up OR remove the self-PID exclusion from a kill helper → `CM-NPROC-HEADROOM-CHECK` FAILs; gate-code = a separate tracked work item).

Canonical authority: constitution submodule `Constitution.md` §12.12. Non-compliance is a release blocker regardless of context. No escape hatch — no `--ignore-thread-limit`, `--dispatch-blind`, `--skip-headroom-check`, `--autonomous-tooling-kill`, `--assume-ulimit-scope` flag exists.

Appendix A – Mutation testing – academic and industrial foundations

The anti-bluff / paired-mutation policy in §1.1 is rooted in the **mutation testing** literature. References that the consuming project SHOULD cite when explaining its meta-test harness:

- Jia, Y. & Harman, M. (2011). *An Analysis and Survey of the Development of Mutation Testing*. IEEE Transactions on Software Engineering, 37(5).
- DeMillo, R.A., Lipton, R.J., Sayward, F.G. (1978). *Hints on Test Data Selection: Help for the Practicing Programmer*. IEEE Computer.
- Open-source mutation testers worth studying:
 - **PIT** (Java) — <https://pitest.org/>
 - **Stryker** (JS / C# / Scala) — <https://stryker-mutator.io/>
 - **Cosmic Ray** (Python) — <https://github.com/sixty-north/cosmic-ray>
 - **mutmut** (Python) — <https://mutmut.readthedocs.io/>
 - **mull** (LLVM-IR) — <https://mull.readthedocs.io/>

The §1.1 paired-mutation policy is a **lightweight** variant of mutation testing applied at the gate-assertion layer rather than at the production-code layer. It does not need a mutation framework; it needs ONE mutation per gate that proves the gate catches the break. This is computationally cheap ($O(\text{gates})$, not $O(\text{production code lines})$) and produces a strong "bluff-immunity" guarantee.

Appendix B – Recursive inheritance & path-independence

Projects that include this constitution submodule MUST tolerate **arbitrary submodule depth**. Between the main project root and any consuming submodule there may be N intermediate levels. To locate this constitution submodule from any depth, every project SHOULD provide a helper that walks up parents until it finds `constitution/Constitution.md` (the file you are reading right now) OR follows `git rev-parse --show-superproject-working-tree` recursively. The constitution submodule ships `find_constitution.sh` for exactly this purpose; nested submodules can source it without knowing their own depth.

§11.4.171 — Mandatory comprehensive human-readable workable-item descriptions (User mandate, 2026-06-29). Every workable item (ATM-NNN ticket) in the project's workable-items database, `Issues.md`, `Fixed.md`, `Issues_Summary.md`, `Fixed_Summary.md`, and ALL derived documents MUST carry a comprehensive human-readable description that explains the item in plain, non-technical language suitable for non-developers (project managers, stakeholders, team leads, business partners). The description MUST be 5-7 sentences minimum, written so that ANY person — regardless of technical background — can understand: (a) WHAT the item is (the feature, fix, or task in simple terms); (b) WHY it matters (the user/business value); (c) HOW it works (the mechanism, without code references); (d) WHO benefits (the end user or stakeholder); (e) WHAT the expected outcome is (the measurable result). The description is MANDATORY — an item without a human-readable description is a §11.4 violation at the documentation layer. Descriptions MUST NOT use jargon, acronyms without expansion, code references, or technical implementation details as the primary explanation. Technical details MAY appear as supplementary context AFTER the plain-language explanation. The `description` column in the SQLite database (`docs/workable_items.db`) is the SINGLE SOURCE OF TRUTH for item descriptions; all derived documents (`Issues.md`, `Fixed.md`, summaries, exports) MUST carry the SAME description text. Pre-build gate `CM-HUMAN-READABLE-DESCRIPTION` (when implemented) walks every item in the DB and asserts the `description` field exists, is non-empty, and is ≥ 50 characters. Paired §1.1 meta-test mutation strips a description → gate FAILs. Classification: universal (§11.4.17). Composes §11.4.15 (item-status tracking), §11.4.16 (item-type tracking), §11.4.19 (column alignment), §11.4.91 (summary clarity), §11.4.93

(SQLite SSoT), §11.4.148 (item integrity), §11.4.65 (universal export). Propagation gate `CM-COVENANT-114-171-PROPAGATION` (literal `11.4.171`) + paired §1.1 mutation.

§11.4.172 — Mandatory production-readiness planning with realistic timeline projection (User mandate, 2026-06-29). Every project under this Constitution MUST maintain a living production-readiness planning document that: (a) provides REALISTIC timeline projections for all product milestones (MVP, beta, production-ready) based on actual development velocity data (items completed per week, reopening rate, blocking dependencies); (b) identifies ALL danger zones, weaknesses, and risk areas with mitigation strategies; (c) accounts for HW delivery timelines, licensing delays, and external dependency risks; (d) defines the critical path and identifies which items can slip without affecting MVP; (e) includes workstation/build-capacity analysis and optimization recommendations; (f) is updated at least monthly or whenever the item count changes by $\geq 10\%$; (g) is cross-referenced against the workable-items database for accuracy. The planning document MUST live at `docs/research/production_planning_<YYYYMMDD>/ANALYSIS.md` with HTML+PDF exports per §11.4.65. Honest boundary (§11.4.6): projections are estimates based on measured velocity — they guarantee the methodology is sound, NOT that specific dates will be met. Classification: universal (§11.4.17). Composes §11.4.6 (no-guessing — projections based on data, not wishful thinking), §11.4.40 (full retest before tag), §11.4.42 (iteration discipline), §11.4.93 (SQLite SSoT), §11.4.108 (four-layer verification). Propagation gate `CM-COVENANT-114-172-PROPAGATION` (literal `11.4.172`) + paired §1.1 mutation.

§11.4.173 — Containerized + distributed build mandate (User mandate, 2026-06-29). EVERY build of EVERY component (source compile, artifact/package/installer/container-image production, codegen, asset render — for ANY language/platform: Go, Android/Gradle, desktop, web, native, firmware) MUST run INSIDE a specialized build container provisioned via the `digital.vasic.containers` submodule (§11.4.76) — NEVER on the bare host. The build containers MUST be DISTRIBUTED to the designated remote build host(s) (e.g. `thinker.local`) via the SAME containers-submodule distribution mechanism the infra uses (§11.4.76 Distributor / remote compose over SSH, §11.4.161 rootless), so the build EXECUTES on the remote build host (offloading the developer/main host); once the build completes the produced artifacts MUST be brought BACK to the originating main host (scp/rsync/volume copy) for use/flashing/distribution. Building outside a container, or on the bare host, is FORBIDDEN — a release blocker (the "works on

my machine" / unreproducible-build class §11.4.76 exists to prevent). The build-host target + build-container definitions are config-injected (§11.4.28, never hardcoded in the submodule); a missing build-container capability is added by EXTENDING the containers submodule upstream (§11.4.74), never by an ad-hoc host build. Honest boundary (§11.4.6): the containerized+distributed build guarantees reproducibility + host-isolation + capacity offload — it does NOT replace the §11.4.40 full-suite retest, §11.4.108 four-layer artifact → runtime verification, or §11.4.38 installable-asset evidence (those still run against the brought-back artifact). Classification: universal (§11.4.17). Composes §11.4.76 (containers submodule — sole orchestration layer), §11.4.161 (rootless runtime), §11.4.74 (extend-don't-reimplement), §11.4.28 (config injection / decoupling), §11.4.24 (build-resource stats), §11.4.82 (iteration speedup — persistent caches in the container), §11.4.121 (no-commit-while-build-writes-artifacts), §11.4.38 (installable-asset evidence on the brought-back artifact), §11.4.108 (artifact → runtime verification), §12.6 (host memory ceiling — offloading the build preserves the main host). Propagation gate `CM-COVENANT-114-173-PROPAGATION` (literal `11.4.173`) + recommended gate `CM-CONTAINERIZED-DISTRIBUTED-BUILD` (every build runs via the containers submodule on the remote build host; a bare-host build is detected + FAILs) + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.173. Non-compliance is a release blocker. No escape hatch — no `--build-on-host`, `--skip-container-build`, `--local-build-ok`, `--no-distributed-build`, `--bare-host-build` flag.

§11.4.174 — Shared-host process-ownership verification mandate (User mandate, 2026-06-29). Verbatim operator mandate: "Make sure that processes we are checking are always OURS! Add this as mandatory constitution rule / constraint, since on every host there may be other work being done!" On ANY host where concurrent work from other projects / users / agents may run (the normal case — a shared developer workstation), EVERY inspection of, conclusion about, or action on a process, build, daemon, port, lock-file, or system-state signal MUST FIRST positively VERIFY the target is OURS (this project's) before drawing a diagnosis or acting on it. A diagnosis based on a process that is actually another project's work is a §11.4.6 guessing-class error; KILLING / signalling / restarting / altering a process that is NOT positively ours is a §12 + §11.4.133-class safety violation (operator-workload destruction) — forbidden. Forensic anchor (FACT, 2026-06-29): during an ATMOSphere build-stall diagnosis, a 200%-CPU `java` + `gradlew` on the shared host were nearly mis-attributed to our SmartTube build;

they were in fact a separate `catalogizer` project's gradle test (`cwd=Projects/catalogizer`), and a naive `pgrep java / gradle` would have matched them as "ours." The mandate (ALL hold): (a) **positive-ownership discriminator**

REQUIRED — attribute a process to us ONLY via a strong signal: its `cwd` resolves inside THIS project's tree, its `argv` names this project's scripts / artifacts / repo path, a PID we ourselves launched and recorded (the §11.4.147 launched-PID registry), or a lock / port / socket path under our tree. A loose name match (`pgrep java / gradle / node / python / podman / zig`) is NEVER sufficient — it matches every project on the host. (b) **NEVER kill / signal / restart a process not positively ours** — other projects' gradle / java / kotlin daemons / podman pods / build processes are strictly off-limits (operator-workload-safety); the §12.8 remediation `pkill -f GradleDaemon` MUST be scoped to OUR daemons (`cwd / argv` under our tree), never a blanket host-wide kill. (c) **contention** → **surface, don't break** — when our work waits on host contention caused by another project's process (e.g. our §12.8 build gate sleeping because another project's gradle daemon is alive), surface it per §11.4.66 / §11.4.101 (the safe blocked-decision) and keep progressing non-blocked work; do NOT kill the other project's process to unblock ours. (d) **filters exclude self + others** — every `pgrep / ps` filter MUST exclude the checking command's OWN `argv` (the self-match that produced the false "still running" restyle watcher bug, 2026-06-29) AND other-project matches; prefer matching on our absolute repo path. (e) **applies to all state, not just processes** — ports, lock files, gradle/maven caches, podman pods, tmux sessions, temp files on a shared host all belong to potentially-other work; verify ownership before reading-as-ours or mutating. Honest boundary (§11.4.6): ownership-verified inspection prevents mis-attribution + cross-project damage — it does not itself prove our process is healthy (that still needs our own artifacts: our build log, our output files, our launched-PID). Composes §11.4.6 (no-guessing) / §11.4.101 (autonomous safe blocked-decision — block-don't-break) / §11.4.66 (surface contention) / §11.4.147 (launched-PID registry as the ownership source) / §12 + §12.8 (host-safety, now scoped to our daemons) / §11.4.133 (don't break other workloads). Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-174-PROPAGATION` (literal `11.4.174`) + recommended gate `CM-PROCESS-OWNERSHIP-VERIFIED` (process/daemon-killing or build-state-diagnosing scripts attribute by `cwd/argv/launched-PID`, never a bare name match; `pgrep` filters exclude self) + paired §1.1 mutation (replace an ownership-verified attribution with a bare `pgrep <name>` → gate FAILs; strip the literal → propagation gate FAILs;

gate-code = separate work item). Non-compliance is a release blocker. No escape hatch — no `--assume-process-ours` , `--loose-pgrep-ok` , `--kill-any-matching-process` , `--skip-ownership-check` flag.

§11.4.176 — Conflict-free multi-track work-division + exactly-once claim registry + capability-aware deadlock-proof device-lock (User mandate, 2026-07-02). When a project runs multiple concurrent work-STREAM tracks (§11.4.167 feature work-stream lifecycle; §11.4.58/§11.4.103 parallel streams) that share (a) a common pool of workable items and (b) a scarce pool of exclusive hardware/resources, TWO cooperating arbitration layers are MANDATORY — both conflict-free BY CONSTRUCTION, not by convention. Forensic anchor — verbatim user mandate (2026-07-02, REM-17): "MUST prevent multiple Tracks working the same workable item OR the same group of workable items resembling one consistent logical group of functionalities. Division of work assigned to Tracks MUST be conflict-free: no conflicts, misunderstandings, race conditions, data corruptions, codebase loss. Risk-free + safe; results ALWAYS produced; no false results, no bluff." (REM-02): "there MUST exist proper synchronization mechanism for flashing and use of device ... lock ... release ... available test devices pool and information about held locks ... dead-lock-proof ... a track locks one OR more devices selected by required capability." **(A) Work-item layer — exactly-once claim registry (REM-17).** Every workable item, OR every logical GROUP of items that share files / subsystem / feature / schema (a hidden-coupling cohort), is assigned to AT MOST ONE track at a time. (1) **Exactly-once** — a lock-arbitrated, append-only claim registry (extending §11.4.147 no-work-loss registry + §11.4.116 real-time sync substrate) records `claim(item|group, track, ts, ttl)` ; a track MUST claim-before-touch; a second claim on an already-claimed item is DENIED (no silent overlap). (2) **Logical-group cohesion** — items that share a file-scope / subsystem / feature / schema are one group claimed WHOLE by one track (prevents the hidden-coupling race two tracks editing the same seam would create — §11.4.58 L3 disjoint-file-scope is the enforcement). (3) **Conflict-free by construction** — disjoint file-scope partition (§11.4.58 L3) + per-track worktree/CoW isolation (§11.4.167) means two tracks physically cannot touch the same bytes. (4) **No codebase loss** — §9.2 hardlinked backups, §11.4.113 ff-only merge-onto-latest-main (NEVER force-push), §11.4.84 working-tree quiescence, §11.4.41 merge-first. (5) **Anti-bluff** — the assignment + the merge are verifiable (registry event stream + captured merge evidence); results are ALWAYS produced (a claimed-but-abandoned item's claim TTL-expires and returns to the pool,

§11.4.147). **(B) Hardware layer — capability-aware deadlock-proof device-lock (REM-02).** When N tracks outnumber M exclusive test devices, exactly ONE track owns each device at a time (§11.4.119 single-resource-owner). A track leases ONE OR MORE devices selected by required CAPABILITY (not merely identity) — the pool records per-device capabilities; leases are capability-matched; multi-device leases are supported. The arbiter is DEADLOCK-FREE by construction — all four Coffman conditions broken: hold-and-wait ELIMINATED (multi-device acquire is ALL-OR-NOTHING inside one critical section — never a partial hold while waiting); circular-wait ELIMINATED (acquire is NON-BLOCKING — wins the whole set at once or returns EBUSY immediately, never waits on another track's device); no-preemption RELAXED (lease TTL + heartbeat + reap — a crashed holder's lease auto-expires, §11.4.147, so a dead track never holds the pool hostage); the only lock ever held is a short bounded flock over the atomic check-and-claim critical section, released before any device work. Registry = append-only JSONL event stream + atomically-rewritten snapshot (§11.4.116). **(C) Universal decoupling + evidence-based resource-tuning (REM-03).** Both layers are universal, decoupled, project-not-aware, reusable (§11.4.17/§11.4.28/§11.4.35) — the consuming project injects its claim-store, device pool, and metric sources at runtime; nothing project-specific lands in the canonical root. Resource-tuning decisions (parallel-stream count, per-stream memory budget §12.6, device-lease breadth, build-queue depth §11.4.167) MUST be EVIDENCE-BASED — driven by captured metrics (§11.4.24 / §11.4.5), never guessed (§11.4.6). Composes §11.4.6 / §11.4.17 / §11.4.24 / §11.4.28 / §11.4.35 / §11.4.41 / §11.4.58 / §11.4.84 / §11.4.85 / §11.4.103 / §11.4.113 / §11.4.116 / §11.4.119 / §11.4.147 / §11.4.167 / §12.6 / §9.2. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-176-PROPAGATION (literal 11.4.176 across the consumer fleet) + recommended gates CM-WORK-DIVISION-EXCLUSIVE-CLAIM + CM-NO-CROSS-TRACK-SCOPE-OVERLAP + CM-LOGICAL-GROUP-COHESION + CM-DEVICE-LOCK-DEADLOCK-FREE (device-lock arbiter ships an all-or-nothing + non-blocking + TTL-reap stress/chaos proof with a paired §1.1 mutation) + paired §1.1 meta-test mutations (gate-code = separate work item). **Canonical authority:** constitution submodule Constitution.md §11.4.176. Non-compliance is a release blocker. No escape hatch — no --allow-cross-track-overlap , --skip-claim-registry , --split-logical-group , --blocking-device-acquire , --allow-partial-device-hold , --tune-resources-by-guess flag.

§11.4.177 — Developer-tooling project-decoupling + invocation-directory operation (User mandate, 2026-07-04). No project-specific script, hook, alias, or binary may be wired — copied, symlinked, PATH-exported, or aliased — into a global or shared developer-tooling PATH; shared tooling **MUST** be project-agnostic and operate on the INVOCATION directory (the cwd, or an explicit path/config argument), **NEVER** on a hardcoded project path. A project-specific helper lives in and runs from its own repo; a genuinely reusable helper is promoted to the shared toolkit **ONLY** after it is stripped of every project-specific assumption (paths, device serials, package names, region endpoints) and takes its target from the invocation context. A project script reachable on the global/shared PATH is a decoupling violation of equal severity to importing project-specific code into a shared submodule (§11.4.28). The forensic case (FACT, 2026-07-04): a project cwd-hook symlinked into the global Claude-Toolkit PATH re-coupled every alias to one hardcoded checkout. Honest boundary (§11.4.6): decoupling guarantees the tool is project-agnostic — it does NOT by itself prove the tool is correct (that still needs the tool's own tests). Composes §11.4.28 (submodules-as-equal-codebase decoupling) / §11.4.29 (naming) / §11.4.35 (canonical-root layer split) / §11.4.11 (file-layout discipline) / §11.4.6 (no-guessing). Classification: universal (§11.4.17) — the consuming project supplies its concrete tool paths + invocation-config mechanism per §11.4.35. Propagation gate `CM-COVENANT-114-177-PROPAGATION` (literal `11.4.177`) + recommended gate `CM-TOOLING-PROJECT-DECOUPLED` (no tracked project script is symlinked/PATH-wired into a global toolkit dir; shared tools take their target from cwd/arg/config, never a hardcoded project path) + paired §1.1 meta-test mutation (wire a project script into the shared PATH, or hardcode a project path into a shared tool → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.177. Non-compliance is a release blocker. No escape hatch — no `--global-symlink-ok`, `--hardcode-project-path`, `--wire-into-shared-path` flag.

§11.4.178 — Track-qualified identity for parallel work streams (session-name collision ban) (User mandate, 2026-07-04). Parallel work streams that share hardware, a git backing store, a tmux/session namespace, a lock namespace, or any single-owner resource **MUST** be addressed by a TRACK-QUALIFIED identity (`<project>__<track>__<role>`), **NEVER** a bare directory basename. Multiple checkouts/clones/worktrees sharing basename `atmosphere` collapse into one session/lock/log key and cross-wire tmux targets, lock files, device claims, and log

directories (the observed collision, 2026-07-04). Every registry row, session name, lock path, log directory, and device claim MUST carry the track-qualified key; a bare-basename identity for a resource with ≥ 2 concurrent claimants is a violation.

Honest boundary (§11.4.6): a unique identity prevents cross-wiring — it does not itself arbitrate ownership (that is §11.4.119 + **§11.4.176** (the multi-track work-division + exactly-once-claim + device-lock arbitration anchor)). Composes §11.4.111 (resolve-by-stable-name-not-by-index) / §11.4.119 (single-resource-owner partitioning) / **§11.4.176** (the multi-track work-division + exactly-once-claim + device-lock arbitration anchor) / §11.4.29 (naming) / §11.4.58 (parallel PWU).

Classification: universal (§11.4.17) — the consuming project supplies its `<project> / <track> / <role>` vocabulary per §11.4.35. Propagation gate `CM-COVENANT-114-178-PROPAGATION` (literal `11.4.178`) + recommended gate `CM-TRACK-QUALIFIED-IDENTITY` (every session name / lock path / log dir / device claim for a multi-claimant resource carries the track-qualified key, never a bare basename) + paired §1.1 meta-test mutation (address a shared resource by bare basename → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item).

Multi-project-per-track layout clarification (2026-07-10): a `/mnt/track<N>` track parent is NOT single-project — MULTIPLE projects coexist on the same track parents, each in its OWN project-named subdirectory `/mnt/track<N>/<project_name>/` (e.g. `/mnt/track1/atmosphere/` AND `/mnt/track1/helix_ota/` side by side); the per-track working copy of project P on track N is ALWAYS `/mnt/track<N>/<project_name>/` (the project's canonical lowercase snake_case dir-name per §11.4.29), never the bare track root and never a shared basename. This is precisely WHY the identity is `<project>__<track>__<role>` and carries the `<project>` component: it keeps two coexisting projects' track-N sessions / locks / logs / device-claims / registry rows from colliding. A project adopting this constitution + the multi-track engine (§11.4.187) MUST create AND register its own `/mnt/track<N>/<project_name>/` per track under its project-qualified identity — it MUST NOT assume it owns the track root or that a track hosts only one project. **Canonical authority:** constitution submodule `Constitution.md` §11.4.178. Non-compliance is a release blocker. No escape hatch — no `--bare-basename-ok`, `--skip-track-qualifier`, `--shared-session-name` flag.

§11.4.179 — Corruption-isolated parallel git streams (own-.git independent clones, not shared-common-dir worktrees) (User mandate, 2026-07-04). When corruption-isolation is a requirement, parallel git work streams MUST each be an

isolated repository with its OWN `.git` (own object store, own index, own lock namespace), NOT `git worktree` checkouts sharing one common `.git`. A shared common-dir is a single point of failure: one stale

`index.lock` / `.commit_all.lock` / `.push_all.lock` freezes EVERY stream at once (the observed 9-h all-track freeze, 2026-07-04), and one corrupted object store loses every stream. Each stream's `git rev-parse --git-common-dir` MUST resolve INSIDE the stream's own tree; a common-dir resolving to a shared parent is a violation. Where CoW/reflink disk-feasibility is the constraint, use the §11.4.167 reflink-clone-with-own-`.git` path — a btrfs/XFS/ZFS/APFS reflink shares extents on disk while keeping a per-stream `.git`, so isolation and disk-feasibility are BOTH satisfied. Honest boundary (§11.4.6): own-`.git` isolation contains lock/corruption blast radius — it does NOT remove the need for per-repo stale-lock reaping (§11.4.180) nor for ff-only no-force integration (§11.4.113). Composes §11.4.167 (feature work-stream lifecycle / CoW clone) / §11.4.119 (single-resource-owner) / **§11.4.176** (the multi-track work-division + exactly-once-claim + device-lock arbitration anchor) / §9.2 (data safety) / §11.4.58 (parallel PWU). Classification: universal (§11.4.17) — the consuming project supplies its reflink mechanism + clone layout per §11.4.35. Propagation gate `CM-COVENANT-114-179-PROPAGATION` (literal `11.4.179`) + recommended gate `CM-GIT-STREAMS-OWN-COMMON-DIR` (every registered parallel stream's `git rev-parse --git-common-dir` resolves inside its own tree, never a shared parent) + paired §1.1 meta-test mutation (register a stream whose common-dir points to a shared parent → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule

`Constitution.md` §11.4.179. Non-compliance is a release blocker. No escape hatch — no `--shared-common-git`, `--worktree-when-isolation-required`, `--skip-git-isolation` flag.

§11.4.180 — Commit/push (single-writer) wrappers MUST auto-reap provably-stale locks before acquiring (User mandate, 2026-07-04). Every commit / push / single-writer wrapper MUST, before acquiring its own lock, auto-reap any PROVABLY-stale git lock — a lock whose recorded holder PID is DEAD (`kill -0` fails), OR (no PID recorded) older than a defined threshold AND with no live wrapper/git process holding that repo. The reap covers the wrapper's own lock plus the git-internal existence-based locks (`index.lock`, `HEAD.lock`, `refs/**/*`) that git does NOT auto-release on holder death. It MUST NEVER remove a lock whose holder is ALIVE (removing a live-held lock corrupts a concurrent writer

— §9.2). Each reap decision (REAPED / KEPT + reason) is logged as captured evidence (§11.4.6 — liveness PROVEN via `kill -0` , never assumed). A wrapper that blocks indefinitely on a dead-holder lock (the observed 9-h freeze, 2026-07-04) is a violation. Honest boundary (§11.4.6): reaping clears dead-holder deadlock — it does NOT substitute for the own- `.git` isolation of §11.4.179 (which shrinks the blast radius from all-streams to one). Composes §11.4.84 (working-tree quiescence) / §11.4.88 (commit-lock-release-immediately + detached push) / §11.4.6 (no-guessing) / §9.2 (data safety) / §11.4.116 (evidence stream) / §11.4.179 (own- `.git` isolation companion). Classification: universal (§11.4.17) — the consuming project supplies its concrete wrapper + reap helper + stale-age threshold per §11.4.35. Propagation gate `CM-COVENANT-114-180-PROPAGATION` (literal `11.4.180`) + recommended gate `CM-WRAPPER-STALE-LOCK-AUTO-REAP` (every single-writer wrapper reaps provably-stale locks before acquiring, never removes a live-held lock, and logs each REAPED/KEPT decision) + paired §1.1 meta-test mutation (make the reaper remove a live-held lock, OR skip the reap-before-acquire step → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.180. Non-compliance is a release blocker. No escape hatch — no `--skip-stale-lock-reap` , `--force-remove-live-lock` , `--block-on-dead-holder` flag.

§11.4.181 — Consistent controlled feature-branch naming: one feature/logic-group ⇒ exactly ONE canonical branch name across the main repo + all owned submodules (User mandate, 2026-07-05). Forensic anchor — verbatim user mandate (2026-07-05): "System MUST pay attention not to create multiple feat. branches with different names for one same feature or group of workable items / features! Naming MUST BE consistent and carefully controlled! This MUST BE part of the constitution too and carefully risk-free handled! No mistakes can happen! There MUST NOT be any false results, mis-alignments, inconsistency or bluff of any kind anywhere!" One feature / logical group of workable items MUST map to EXACTLY ONE canonical branch name, used IDENTICALLY on the main repo AND on every owned submodule the group touches — NEVER two (or more) differently-named branches for the same feature/group, NEVER a per-repo naming drift. The mandate (ALL hold): (1) **Single canonical name per group.** For each feature/logical-group a single canonical branch name is minted ONCE (per the §11.4.167 D scheme: `feature/<slug>` on the branch axis + `<prefix>-<base>-feat-<slug>` on the tag axis, `<slug>` lowercase snake/kebab per §11.4.29). That

exact string is the branch name on the main repo AND on every touched owned submodule; an untouched submodule stays on its base branch (it is NOT force-branched), but any submodule that IS branched for this group MUST use the identical canonical name. (2) **Single source of truth for branch names — look up, never re-invent.** The canonical name is recorded ONCE in a tracked registry and every subsequent branch creation for that group LOOKS IT UP rather than re-deriving/re-typing it. Bind this to the §11.4.176 exactly-once claim registry / the §11.4.93 workable-items DB `logic_group → branch` mapping (a `branch_name` field on the group/claim record): the claim that reserves a feature/group ALSO reserves its canonical branch name, so a second stream working the same group reads the same name instead of minting a divergent one. Creating a branch whose name does NOT match the registry's canonical name for that group — OR minting a SECOND canonical name for a group that already has one — is a §11.4.181 violation (a §11.4.6 no-guessing violation: the branch name is a looked-up FACT, never an ad-hoc re-invention). (3) **Mechanical enforcement (§11.4.75).** A pre-build/commit gate `CM-BRANCH-NAME-CONSISTENCY` verifies that every live feature branch (main repo + each owned submodule) maps to a registered feature/logical-group AND carries the identical registered canonical name across the main repo and every touched submodule; it FAILs on (a) an unregistered feature-branch name, (b) two differently-named branches bound to the same group, (c) a submodule branch whose name diverges from the group's main-repo canonical name. Paired §1.1 mutation: register a single group under two divergent branch names (or point a submodule branch at a name ≠ the group's canonical name) → the gate FAILs; restore → PASS. (4) **Risk-free reconciliation — no force, no loss (§9.2 / §11.4.113 / §11.4.84).** When a mis-named / duplicate branch is discovered, reconciling it to the canonical name MUST preserve every commit: create/checkout the canonical branch at the merge-base, MERGE the mis-named branch's commits onto it (union, no `-s ours` /rebase/reset that drops commits per §11.4.113 step 3), verify no commit lost (`git log <canonical>..
<misnamed>` empty after merge) + no conflict markers, push fast-forward to all upstreams (NEVER force-push — the canonical branch descends from the mis-named tip), THEN retire the mis-named branch only after the merge is confirmed on all mirrors; delete-before-merge is forbidden. Honest boundary (§11.4.6): "the branches are probably the same feature" is NOT a determination — bind the branch to a group via the registry (claim id / logic-group id), never guess two branches are "the same feature" without the registry record. Classification: universal (§11.4.17) — a platform-neutral naming-discipline reusable across any multi-repo/multi-submodule project; the consuming

project supplies its concrete registry (claim registry / workable-items DB), branch/tag scheme, and owned-submodule set per §11.4.35. Composes §11.4.167 (feature work-stream lifecycle / D branch+tag scheme) / §11.4.176 (exactly-once claim registry — the branch-name reservation rides the claim) / §11.4.178 (track-qualified identity — a track-qualified session/lock key is distinct from the shared cross-track feature-branch name; both hold) / §11.4.179 (own- `.git` streams still share ONE canonical feature-branch name across their independent repos) / §11.4.28 (owned-submodule set) / §11.4.29 (lowercase snake/kebab slug) / §11.4.93 (workable-items DB `logic_group`) / §11.4.113 (no-force merge-onto-latest) / §11.4.84 (quiescence) / §9.2 (no-loss) / §11.4.6 (branch name is a looked-up fact) / §11.4.75 (mechanical enforcement). Propagation gate

`CM-COVENANT-114-181-PROPAGATION` (literal `11.4.181` across the consumer fleet) + recommended gate `CM-BRANCH-NAME-CONSISTENCY` + paired §1.1 meta-test mutation (register one group under two divergent branch names → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.181. Non-compliance is a release blocker. No escape hatch — no `--allow-divergent-branch-names`, `--skip-branch-registry`, `--rename-branch-freely`, `--per-repo-branch-name-OK`, `--two-names-one-feature-OK` flag.

§11.4.182 — Track+branch work-stream identity label (`T<N>/<branch>`) on every agent/subagent/work-stream label and every operator-facing work-stream reference (User mandate, 2026-07-05). Forensic anchor — operator mandate (2026-07-05): "EVERY agent/subagent/work-stream label AND every operator-facing report reference to a work stream MUST carry a

`(T<N>/<branch>)` prefix identifying the track + git branch — e.g. `(T1/main)` `ATM-312 MPV drm_prime investigation`, `(T2/feat/mistiq-vader)` `SPK-XXX ...`." In a multi-track layout (parallel checkouts each on a distinct branch — e.g. `/mnt/track1..4`) a bare work-stream label ("ATM-312 investigation") does not say WHICH track/branch produced it, so cross-track dispatches, status reports, and progress lines collide and mislead. The mandate (ALL hold): (1) **Universal label form.** Every agent / subagent / work-stream label AND every operator-facing reference to a work stream (agent-dispatch descriptions, status reports, resume/handoff docs, progress lines) MUST START with a `(T<N>/<branch> - <alias>)` prefix — `<N>` the track number, `<branch>` the git branch, `<alias>` the claude-toolkit alias in use — e.g. `(T1/main - claude1)` `ATM-312 MPV drm_prime investigation`, `(T2/feature/mistiq-vader - deepseek)` `SPK-XXX subtitle`

work . (2) **Deterministic derivation, never guessed (§11.4.6).** The label is DERIVED, not typed from memory: `<N>` from the checkout path (the multi-track root `/mnt/track<N>/...` ; honest `?` when the checkout is not under a numbered track — never a guessed number), `<branch>` from `git rev-parse --abbrev-ref HEAD` , `<alias>` from `CLAUDE_CONFIG_DIR` `basename .claude-<alias>` → `<alias>` (honest `?` when unset/non-matching, never a guessed name). The consuming project supplies a deterministic labeler (reference impl `scripts/multitrack/track_branch_label.sh`). (3) **Mechanical enforcement (§11.4.75 / §11.4.109-class anti-forgetting).** A PreToolUse guard hook (reference impl `constitution/scripts/hooks/guard-track-branch-label.sh` , inherited BY REFERENCE per §11.4.177, never copied) BLOCKS at the tool-call boundary any agent-dispatch tool (Agent / Task / TaskCreate) whose `description` (or `subagent` `label`) does not START with `^(T[0-9]+/[^\]]+\)` — so enforcement is independent of any agent's recall (a rule the orchestrator forgets to paste is not enforcement); non-agent tools pass through untouched. A pre-build gate `CM-TRACK-BRANCH-LABEL` asserts the labeler + hook exist, are executable + parseable, the hook is wired in the runtime settings, and the convention doc exists; a propagation gate `CM-COVENANT-114-182-PROPAGATION` asserts the literal `11.4.182` is present across the governance-file fleet. Paired §1.1 mutation: strip the hook's prefix-check → the hook test-suite FAILs; strip the `11.4.182` literal from a governance file → the propagation gate FAILs; restore → PASS. (4) **Honest boundary (§11.4.6).** A `?` in the track OR the alias field is a valid honest label (unknown track / no git repo), never a bluff; the enforced FORMAT is numeric-track `[0-9]+` per the reference regex, so an agent dispatch from a non-track checkout is surfaced (blocked) rather than silently mislabeled. Classification: universal (§11.4.17) — a platform-neutral multi-stream identity discipline reusable by any project running parallel track/branch work streams; the consuming project supplies its concrete track-root layout, labeler, hook wiring, and governance-file set per §11.4.35. Composes §11.4.178 (track-qualified identity — §11.4.182 is the label-prefix instantiation applied to EVERY agent/subagent label + every operator-facing work-stream reference) / §11.4.29 (naming) / §11.4.75 (mechanical enforcement) / §11.4.109 (anti-forgetting PreToolUse guard) / §11.4.177 (hook inherited by reference, project-decoupled) / §11.4.6 (label is a derived fact; `?` is honest-unknown) / §11.4.58 · §11.4.103 (parallel work streams) / §11.4.119 (single-resource-owner per track) / §11.4.116 · §11.4.147 (work-stream registry identity). Propagation gate `CM-COVENANT-114-182-PROPAGATION` (literal `11.4.182` across

the consumer fleet) + recommended gate `CM-TRACK-BRANCH-LABEL` + paired §1.1 meta-test mutation. **TRUNK RULE (operator mandate, 2026-07-28):** work performed on the canonical trunk (`main / master`) is **ALWAYS Track 1** — the label reads (`T1/main - ...`) , NEVER (`T?/main - ...`) , regardless of checkout path. Precedence is deliberate: the BRANCH decides first and the checkout path decides only for non-trunk branches, so a trunk checkout living outside `/mnt/track<N>` — or inside a differently-numbered one — is still Track 1; the identity that matters is which STREAM the work belongs to, not which directory it happens to sit in. Rationale: §11.4.182 labels are how concurrent streams are told apart, so an unnumbered `T?` on trunk work makes it indistinguishable from a genuinely-unknown track in every dispatch, report and progress line — and the trunk is a well-known, singular stream that is never unknown. Reference implementation:

`scripts/multitrack/track_branch_label.sh` derives the branch BEFORE the track number and maps `main|master -> 1` . Labels MUST be re-derived from that script at dispatch time and kept in real-time sync with the live branch/alias/model/effort — a remembered or stale label is itself a §11.4.182 violation (§11.4.6: live identity is not knowable from memory). **Canonical authority:** constitution submodule `Constitution.md` §11.4.182. Non-compliance is a release blocker. No escape hatch — no `--skip-track-label` , `--unlabeled-agent-OK` , `--bare-work-stream-reference` , `--no-track-branch-prefix` , `--label-from-memory` flag.

§11.4.183 — Maximal multi-agent utilization per work-stream + full-constitution-application + zero-bluff mandate (User mandate, 2026-07-07).

Forensic anchor — verbatim operator mandate (2026-07-07): "Every Track MUST USE and maximize multi-agents working approaches (subagent-driven, independent agents for various reviewing, and other forms of agents use)! We MUST maximize use of agents so every track produces as much as possible completed work with highest possible quality for the least spent time! Every Track MUST apply everything from the constitution — all rules, mandatory constraints, guides and technology derived! Nothing can be ignored, skipped, avoided or considered less relevant! There MUST NEVER BE false or faulty results, untested and unverified work or bluff of any kind anywhere!" Every track / work-stream MUST maximize multi-agent working approaches — subagent-driven development (§11.4.20/§11.4.70), independent review agents (§11.4.125/§11.4.142/§11.4.165/§11.4.134), parallel background streams with auto-backfill (§11.4.58/§11.4.103), and every other applicable agent form — so that every track produces the MOST

completed work at the HIGHEST quality in the LEAST time. The mandate (ALL hold): (1) **Maximize agent utilization per track.** Every track/work-stream (main + every parallel background stream per §11.4.58/§11.4.103) MUST default to subagent-driven execution for non-trivial work (§11.4.20/§11.4.70), dispatch independent review agents structurally separated from the author for quality/safety/correctness review (§11.4.125/§11.4.142/§11.4.165, iterated to a clean GO per §11.4.134), and use every other applicable specialised-agent from the runtime offers — never one monolithic foreground agent doing everything a task naturally decomposes into. (2) **Full-constitution application — nothing ignored.** Every track MUST apply the ENTIRE constitution to its work — every universal rule, every mandatory constraint, every guide, and every derived technology (gates, mutation pairs, sync engines, recording/evidence pipelines, lock/claim registries) this Constitution and its consuming project's governance define. No rule may be silently ignored, skipped, avoided, or judged "less relevant" for a given track — partial-constitution-application is itself a §11.4.183 violation, equal in severity to skipping a mandated gate. (3) **Zero-bluff, zero-false-result, zero-untested-work.** There MUST NEVER be false or faulty results, untested or unverified work, or bluff of any kind, anywhere, on any track — this restates and BINDS the entire §11.4 anti-bluff covenant (§11.4/.1/.2/.5/.6/.7/.50/.69/.107) as an unconditional per-track invariant, not merely a release-gate check: every claim of "done" on any track carries captured evidence (§11.4.5/§11.4.69/§11.4.107), every gate/mutation pair is genuinely falsifiable (§1.1), and no track may report progress it did not actually perform (§11.4.126). Honest boundary (§11.4.6): "maximize agents" is bounded — by §12.6's 60% host-memory ceiling, §11.4.58's parallel-agent cap, genuine task parallelizability, and §11.4.119's single-resource-owner rule for contended hardware/locks — so spawning contending, redundant, or non-parallelizable agents is NOT maximization; the bar is maximum USEFUL parallel throughput per track, never raw agent count. Classification: universal (§11.4.17) — a platform-neutral operating-discipline reusable by any project running one or more parallel work-streams/tracks; the consuming project supplies its concrete track layout, subagent runtime, and review-agent roster per §11.4.35. Composes §11.4.20 (subagent-driven-by-default) / §11.4.58 (parallel-development PWU pipeline) / §11.4.70 (subagent-driven execution is the default) / §11.4.92 (multi-pass change evaluation) / §11.4.94 (zero-idle priority-first parallel-by-default) / §11.4.97 (maximum-use-of-idle-time) / §11.4.103 (continuous parallel-stream working routine) / §11.4.119 (single-resource-owner partitioning) / §11.4.125 (code-review-agent gate) / §11.4.126 (default autonomous-loop working mode) / §11.4.134 (code-

review iterate-until-GO) / §11.4.142 (universal code-review mandate) / §11.4.165 (universal independent verification agent mandate) / §12.6 (60% memory ceiling). Propagation gate `CM-COVENANT-114-183-PROPAGATION` (literal `11.4.183` across the consumer fleet) + recommended gate `CM-MAX-AGENTS-PER-TRACK` (every active track shows subagent/review-agent dispatch evidence for non-trivial work, not solitary foreground execution) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; record a track that skipped subagent/review-agent dispatch for non-trivial work with no genuine single-agent justification → `CM-MAX-AGENTS-PER-TRACK` FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.183. Non-compliance is a release blocker. No escape hatch — no `--single-agent-OK` , `--skip-review-agents` , `--partial-constitution-application` , `--serialise-parallelisable-work` flag.

§11.4.184 — Mandatory SonarQube static-analysis CLI + local tooling installed and PATH-discoverable (User mandate, 2026-07-06). Forensic anchor — operator mandate (2026-07-06): "Make sure sonarqube is installed and configured locally same we did for Firebase, GitHub and GitLab CLIs through exported System PATH and .bashrc! Add all major points to the constitution so all projects do the same!" Every project/host that runs static analysis MUST have the SonarQube scanner CLI (`sonar-scanner`) installed AND durably discoverable on the system PATH via the shell rc (`.bashrc` / `.zshrc`) — exactly like the Firebase/GitHub/GitLab CLIs — PLUS the shared `constitution/scripts/sonarqube/tooling` (`sonarqube_install_check.sh` / `sonarqube_lib.sh` / `sonarqube_run_scan.sh` / `sonarqube_container.sh` + `compose/`) consumed **by reference** (§11.4.28 decoupled + §11.4.177 project-agnostic + §11.4.80-style inherited, NEVER copied per project). The mandate (ALL hold): (1) **Installed + PATH-durable** — the scanner CLI resolves in every fresh shell (a durable rc export — e.g. a generic `Factory/software/*/bin` glob or an explicit line — not an ephemeral session-only PATH); (2) **Shared tooling present + GREEN** — `sonarqube_install_check.sh` exits 0 (scanner + rootless podman per §11.4.161 + podman-compose + curl + ES-ready host `vm.max_map_count`) before any scan-dependent work; (3) **Anti-bluff (§11.4.6)** — "installed/configured" is PROVEN by `command -v sonar-scanner` resolving AND the install-check exit 0, NEVER assumed; a claimed-installed-but-absent CLI is a §11.4 bluff at the tooling layer; (4) **Rootless (§11.4.161)** — the local SonarQube server runs via rootless podman through the container tooling, never rootful docker/sudo. **Honest boundary**

(§11.4.6): this anchor names the SonarQube SCANNER CLI + the shared local-server tooling; per §11.4.166 the former universal Semgrep static-analysis mandate is REPEALED — SonarQube is a DISTINCT, operator-mandated tool here, not a Semgrep re-instatement, and this anchor mandates the TOOLING availability, not that every project must run a scan on every build. Classification: universal (§11.4.17) — the consuming project/host supplies its concrete install path + rc-export mechanism per §11.4.35. Composes §11.4.28 (decoupled shared tooling) / §11.4.36 (install-on-clone) / §11.4.74 (submodule-catalogue-first) / §11.4.75 (mechanical enforcement) / §11.4.109 (anti-forgetting) / §11.4.161 (rootless container) / §11.4.166 (Semgrep-repeal boundary) / §11.4.177 (project-decoupled tooling) / §11.4.6 (no-guessing). Propagation gate `CM-COVENANT-114-184-PROPAGATION` (literal `11.4.184` across the consumer fleet) + recommended gate `CM-SONARQUBE-CLI-INSTALLED` (`command -v sonar-scanner` resolves + `sonarqube_install_check.sh` exit 0 + shared tooling present) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; move the scanner off PATH → the install gate FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.184. Non-compliance is a release blocker. No escape hatch — no `--skip-sonarqube-install`, `--sonar-cli-optional`, `--ad-hoc-sonar-path`, `--rootful-sonar-OK` flag.

§11.4.185 — Manual QA-team testing as the mandatory FINAL confirmation of done (User mandate, 2026-07-07). Forensic anchor — verbatim operator mandate (2026-07-07): "We MUST EVERYWHERE and ALWAYS take into account that manual testing as final confirmation step for work done is MANDATORY!!! Without such confirmation by QA team of ours we cannot consider any set of features or scope of work fully completed! This MUST BE well known and mandatory part of our constitution! Fully respected, taken into account, followed and never forgotten!" No set of features / scope of work / release may be considered **fully completed** until it has received MANUAL TESTING confirmation by the project's QA team as the FINAL confirmation step — EVERYWHERE and ALWAYS (every feature, every scope, every release), NEVER forgotten. The mandate (ALL hold): (1) **Necessary-but-not-sufficient automation.** Every automated gate this Constitution mandates — the §11.4 anti-bluff covenant, §11.4.5/§11.4.69 captured evidence, §11.4.52 autonomous validation, §11.4.40 full-suite retest, §11.4.108 four-layer SOURCE → ARTIFACT → RUNTIME → USER-VISIBLE verification, §11.4.132 risk-ordered validation — is NECESSARY but NOT SUFFICIENT for "fully completed." These are the automated confidence gates; the

QA-team MANUAL test is the FINAL human sufficiency gate that closes the scope.

(2) **Composition with §11.4.52 — ADDS, does not weaken.** §11.4.52 requires every feature to have an AUTONOMOUS validation path (primary, CI-scalable, runs on every commit, survives operator absence) and forbids operator-attended-ONLY validation. §11.4.185 does NOT weaken that requirement — it ADDS a final human QA sign-off ON TOP of the mandatory autonomous path. A feature MUST have BOTH: the autonomous path (§11.4.52, primary + always-on) AND the manual-QA final confirmation (§11.4.185, the closing gate). Manual QA is the FINAL gate, never the ONLY gate — a feature validated ONLY by manual QA with no autonomous path is STILL a §11.4.52 violation; a feature validated ONLY by automation with no manual-QA sign-off is a §11.4.185 violation. Both must hold. (3) **Workflow placement.** The delivery pipeline is: autonomous-validation-GREEN → build → flash/deploy → autonomous on-device validation → **HAND OFF to the QA team for manual testing** → ONLY after QA-team manual confirmation is the scope "fully completed" / eligible for closure or a release tag. The §11.4.126 autonomous-loop terminal condition (a fully-validated-and-verified release tag) now REQUIRES the QA-team manual confirmation before the tag is cut; a scope marked "done" / "fully completed" / migrated to Fixed.md without a recorded QA-team manual confirmation is a §11.4.185 violation — a definition-of-done bluff of the same severity class as a §11.4 PASS-bluff. (4) **Honest boundary (§11.4.6).** The agent/loop MAY autonomously drive the scope through autonomous-GREEN + build + flash + on-device autonomous validation, but it MUST then hand off + WAIT for the QA-team's manual confirmation before claiming "fully completed" — it MUST NOT self-certify the manual step, simulate it, or infer it from automated evidence alone (that would itself be a bluff at the human-sufficiency layer). While awaiting QA-team confirmation the agent keeps progressing every OTHER non-blocked item in parallel (§11.4.87/§11.4.94/§11.4.97/§11.4.101) — the wait parks that one scope, not the loop. Classification: universal (§11.4.17) — a platform-neutral definition-of-done discipline reusable by ANY project that maintains a QA team; the consuming project supplies its concrete QA-team identity, hand-off mechanism, and confirmation-recording location per §11.4.35. Composes with §11.4 (anti-bluff covenant) / §11.4.5 (captured-evidence quality) / §11.4.6 (no-guessing — the manual confirmation is a recorded FACT, never assumed) / §11.4.40 (full-suite retest before tag) / §11.4.52 (autonomous-validation mandate — §11.4.185 is its human-sufficiency CAPSTONE, adding rather than substituting) / §11.4.69 (sink-side evidence taxonomy) / §11.4.108 (four-layer fix-verification) / §11.4.126 (autonomous-loop terminal condition now gated on QA-team confirmation) /

§11.4.129 (huge-blocker release protocol) / §11.4.130 (post-remediation validate-the-fix-first) / §11.4.132 (risk-ordered validation priority). Propagation gate `CM-COVENANT-114-185-PROPAGATION` (literal `11.4.185` across the consumer fleet) + recommended gate `CM-MANUAL-QA-FINAL-CONFIRMATION` (every item marked fully-completed / migrated to Fixed.md carries a recorded QA-team manual confirmation citation) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; mark an item fully-completed with no QA-team confirmation citation → the recommended gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.185. Non-compliance is a release blocker regardless of context. No escape hatch — no `--skip-manual-qa`, `--automation-suffices`, `--self-certify-manual-step`, `--assume-qa-confirmed`, `--autonomous-only-completion` flag.

§11.4.186 — Anti-divergence enforcement: cross-document consistency is a mandatory-before-export/commit gate, never an after-the-fact audit (research-derived, 2026-07-08). Compact summary: any project maintaining more than one representation of the same tracked data (a delivery plan + its MVP brief + a workable-items DB + their summaries + `.md` / `.html` / `.pdf` / `.docx` exports) MUST enforce cross-document CONSISTENCY as a deterministic gate that runs BEFORE any export render, any doc/DB sync `verify`, and any doc-set commit — NEVER as an after-the-fact human audit that runs only when an operator happens to ask (the forensic FACT: the same NanoKVM/phone-as-screen feature appeared in three plan clusters — one shared ticket SPK-596 across SIX rows split over TWO releases, and a carve-out due to ship 2026-08-31 whose own prerequisite does not START until 2026-10-05 — and NO gate caught it; only an operator's manual read of `Plan_v9.0.xlsx` surfaced it, after the divergent bytes had already exported). (1) **One no-divergence verdict, three seams** — a single deterministic PASS/FAIL/SKIP verdict (built by composing a cross-document integrity validator with the workable-items-DB internal `validate`) gates all three write-seams (export / sync-verify / commit); a FAIL at any seam refuses that seam, naming the exact offending records (§11.4.6 actionable). (2) **Five decidable check families** — DEDUP (no two records describe one feature with divergent timeline/status/release, keyed on ticket OR normalised `(subject, scope)`, NOT bare subject substring), TIMELINE (start ≤ deadline; no deadline-before-dependency; no dependency cycle; GATED exempt-but-marked), CROSS-DOC (the same item across every representation carries identical timeline/status/type against a NAMED authoritative

source), INTEGRITY (no orphan refs; Status ↔ Type §11.4.33; location ↔ status), STRUCTURAL (required columns; ID uniqueness+monotonicity §11.4.54). (3) **Drift-proof trigger** — a §11.4.86 sha256-of-sorted-members fingerprint of the authoritative inputs re-arms the gate on ANY input change, so a stale verify cannot pass on changed data. (4) **Self-validated analyzer (§11.4.107(10))** — the validator ships golden-good (MUST PASS) + one golden-bad per family (MUST FAIL, pinpointing the offender) + a negative-control (distinct same-subject tasks that MUST PASS — the false-positive guard); a validator that passes its golden-bad, or fails golden-good/the negative-control, is itself a §11.4 bluff and a release blocker. (5) **Reusable + decoupled (§11.4.28/.31)** — packaged as a depth-1 reusable engine under the constitution submodule per the §11.4.28(C) carve-out with a `helix-deps.yaml` ; project-specific doc-sets + authoritative-source bindings live in a consumer-owned checkset (data, never engine code). (6) **Supersedes ad-hoc audits** — the per-cycle `*_INTEGRITY_FINDINGS.md` / `RECONCILIATION_*.md` after-the-fact audit docs are retired in favour of the gate; honest boundary §11.4.6 — the gate proves internal CONSISTENCY (no record diverges, no timeline is incoherent, every ref resolves), NOT plan CORRECTNESS (achievability/date-rightness stay operator/management decisions the gate SURFACES, never MAKES). Classification: universal (§11.4.17). Composes §11.4.6/.12/.33/.44/.50/.53/.54/.60/.65/.73/.75/.85/.86/.91/.93/.95/.106/.107/.108/.110/.148/.176. Propagation gate `CM-COVENANT-114-186-PROPAGATION` (literal `11.4.186`) + recommended gate `CM-DOC-INTEGRITY-VALIDATION` (5 invariants: validator present+executable; consumer checkset present; the no-divergence gate wired into each of the three seams; fingerprint sidecar fresh; `selfcheck` golden-good/golden-bad/negative-control wired into meta-test) + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.186. Non-compliance is a release blocker. No escape hatch — no `--skip-divergence-gate` , `--audit-after-export-OK` , `--cross-doc-check-not-applicable` , `--unvalidated-analyzer-OK` , `--divergent-commit-OK` flag.

§11.4.187 — Automatic multi-track ruler orchestration (User mandate, 2026-07-04; landed 2026-07-09). Compact summary: every project incorporating this constitution MUST ship its multi-track parallel-development orchestration as a UNIVERSAL, AUTOMATIC, OUT-OF-THE-BOX, INHERITED capability — not a per-project reimplementations — via a single conductor session that programmatically SPAWNS, DRIVES, RESUMES, MONITORS, and CONTROLS

per-track headless workers. (1) **Headless spawn primitive** — workers launch via native headless `claude -p --output-format stream-json --verbose`; the worker's `session_id` is captured from the FIRST `init` event (never invented); success/failure is read from the terminal event's `result.is_error` field, NEVER the process exit code (a CLI can exit non-zero on a successful `result`, or zero on a failed one — exit-code-as-success-signal is a §11.4.6 guess); continuation of an existing worker uses `--resume <session_id>` invoked from the SAME cwd + env the spawn used (the storage/resume-collision gotcha). (2) **Per-subscription auth + rate-limit fallback** — every spawn `unset` s any ambient `ANTHROPIC_API_KEY` (the precedence trap that silently routes a "subscription" worker onto pay-per-token API billing) and sets the selected alias's `CLAUDE_CODE_OAUTH_TOKEN` / `CLAUDE_CONFIG_DIR`; on a captured rate-limit signature (`isApiErrorMessage + apiErrorStatus:429`, quota-exhausted vs transient distinguished by evidence, never guessed) the ruler REBINDS the affected track onto the next healthy alias (cross-track reuse permitted), falls back to a provider alias where configured, and bounded-parks (§11.4.101) rather than hanging when every alias is cooled. (3) **Crash-resilient ruler self-supervisor** — the ruler's OWN state (per-track alias/session_id/worktree/last-verdict) is persisted durably (write-temp-then-rename, NEVER tmpfs) so an external watchdog detects a dead ruler PID and re-hydrates every track's binding from the durable state + the §11.4.116 event stream — no work is silently lost to a ruler crash (§11.4.147 at the ruler layer). (4) **Idempotent bootstrap + auto-install + conductor-stays-home** — one idempotent bootstrap command wires the whole system out of the box (cwd-hook install, config validation, reconciliation), auto-invoked on every constitution pull via the §11.4.164 `post_update_hook.sh` skill-register seam; the `conductor`: -designated alias/session is NEVER bound into a track worktree (the resolver returns nothing for it) so the conductor always runs from the main checkout. (5) **Project-agnostic engine, consumer-owned data** — the entire mechanism (config loader, alias → worktree resolver, cwd-hook, bind/fallback/cooldown orchestrator, session launcher, rate-limit monitor, supervisor) lives at `constitution/scripts/multitrack/`, inherited BY REFERENCE (§11.4.28(B)/§11.4.177), carrying NO project literal; every consuming project supplies its own per-host `config/multitrack/<hostname>.yaml` (tracks/mounts/serials/aliases/conductor) as DATA, never as an engine edit. Anti-bluff (§11.4/§11.4.108): every piece carries four-layer coverage (pre-build gate + on-host test + paired §1.1 mutation + captured evidence per §11.4.69), and each fix declares a runtime

signature verified on a clean host — never a claim the mechanism "should work."

Honest boundary (§11.4.6): this anchor mandates the MECHANISM; genuine per-subscription quota-isolation proof requires live tokens exercising real rate-limit boundaries (an operator-scheduled live acceptance window) — a mock-driven GREEN proves the spawn/resume/fallback CONTRACT, never live-quota behaviour, and the two are never conflated or substituted for each other.

Classification: universal (§11.4.17). Composes §11.4.20 (subagent-driven-by-default) / §11.4.28 (submodule decoupling + dependency layout) / §11.4.58 (parallel-development PWU pipeline) / §11.4.70 (subagent-driven execution default) / §11.4.101 (autonomous-decision-over-blocking — the bounded-park-on-all-cooled behaviour) / §11.4.103 (continuous parallel-stream routine) / §11.4.111 (resolve-by-stable-name — config by serial, never by index) / §11.4.116 (real-time conductor ↔ framework sync channel) / §11.4.119 (single-resource-owner partitioning) / §11.4.126 (default autonomous-loop mode) / §11.4.128 (always-on device-recording — the analogous host-safety discipline) / §11.4.131 (standing session-resumption file) / §11.4.147 (crashed-agent respawn-until-complete — the ruler-self-supervisor is its ruler-layer instantiation) / §11.4.164 (constitution auto-propagation hook seam) / §11.4.167 (feature work-stream lifecycle) / §11.4.176 (multi-track work-division + exactly-once claim + device-lock) / §12.6/§12.8 (host-safety memory + concurrency ceilings — the spawn/launch pool is budget-gated).

Propagation gate `CM-COVENANT-114-187-PROPAGATION` (literal `11.4.187`) + recommended gates `CM-MULTITRACK-ENGINE-IN-CONSTITUTION` (the generic ruler scripts present under `constitution/scripts/multitrack/`, carrying no project literal, `bash -n + sh -n` clean per §11.4.67) / `CM-MULTITRACK-AUTOINSTALL-WIRED` / `CM-MULTITRACK-BOOTSTRAP-IDEMPOTENT` / `CM-MULTITRACK-BIND-ON-START` / `CM-MULTITRACK-FALLBACK-MONITOR` / `CM-MULTITRACK-HEADLESS-SPAWN` / `CM-MULTITRACK-RULER-SUPERVISOR` + paired §1.1 mutation (strip the literal → propagation gate FAILs; re-introduce a project literal into the engine, or downgrade a rate-limit rebind to exit-code-based success detection → the corresponding mechanism gate FAILs). **Canonical authority:** constitution submodule `Constitution.md` §11.4.187. Non-compliance is a release blocker. No escape hatch — no `--skip-multitrack`, `--manual-multitrack-only`, `--no-auto-install`, `--per-project-reimpl-OK`, `--exit-code-is-success-signal` flag.

§11.4.188 — Regular main → feature merge cadence: every track / feature-branch MUST FREQUENTLY merge canonical `main` into itself DURING the work, never only at the end (User mandate, 2026-07-09). Compact summary: every long-lived feature branch AND every parallel-development track (§11.4.58/§11.4.103 parallel streams; §11.4.167 feature work-streams; §11.4.176/§11.4.178/§11.4.179 multi-track) MUST regularly `git merge origin/main` (the canonical trunk/master) INTO its own branch THROUGHOUT the work — NOT only at the end — so no branch ever drifts far from main and the eventual back-merge stays a small, low-conflict, low-risk operation. GENERALISES §11.4.167(D) (trunk-merged-INTO-every-BIG-feature-STREAM, at minimum after every trunk tag / daily, `git merge` never rebase, stale-if-not-synced) from the big-work-item feature-stream lifecycle to EVERY feature branch on EVERY track — a small/medium feature branch, or any `git worktree` /CoW track that never earned a full §11.4.167 stream, is STILL bound to this cadence (the §11.4.167(D) discipline promoted to the general standalone rule; §11.4.167(D) remains its big-item specialisation). (1) **WHAT** — fetch the latest canonical `main` / `master` and merge it INTO the working branch, keeping the branch continuously close to trunk; a branch that has NOT integrated trunk within the cadence is flagged STALE (the exact long-lived-branch conflict-storm + integration-risk anti-pattern this forbids). (2) **CADENCE** — merge trunk in FREQUENTLY: after EVERY trunk tag, at minimum DAILY while the branch is live, AND before starting any significant new chunk of work on the branch — MANY SMALL merges, never one big deferred end-of-branch merge (small frequent merges minimise the divergence + conflict surface). (3) **HOW (safe-by-construction)** — MERGE, NEVER rebase a shared/tagged/pushed branch (no commit-loss, no history rewrite — §9/§11.4.113); FETCH-FIRST every remote before the merge (§11.4.37/§11.4.71 — remote state is unknowable without a fetch, §11.4.6); take a §9.2 hardlinked pre-op backup before any LARGE / conflict-heavy / risky merge; resolve every conflict CAREFULLY — ZERO conflict markers committed, ZERO file silently dropped, union of both sides preserved (§11.4.41 step 3 / §9 no-loss); integrate ONLY when the branch is QUIESCENT (§11.4.84 — no in-flight mutation gate, no unaccounted uncommitted work); prefer running the merge + its verification in the BACKGROUND parallel to the main stream so it never stalls other work (§11.4.88/§11.4.89/§11.4.103); NEVER force-push the merged result (§11.4.113 — the merged branch fast-forwards its own mirrors). (4) **CONSISTENCY + ANTI-BLUFF** — after EVERY merge the branch MUST be PROVEN consistent, never ASSUMED: a pre-build smoke gate runs post-merge (§11.4.42 step-3 smoke) and is GREEN, a `git grep` for conflict markers

(<<<<<< / ===== / >>>>>>) returns EMPTY, and no commit present pre-merge on either side is lost — captured evidence per §11.4.5/§11.4.69; a merge claimed "clean" without the post-merge smoke + marker-scan + no-lost-commit check is a §11.4/§11.4.1 bluff at the integration layer. (5) **HONEST BOUNDARY (§11.4.6)** — regular trunk-integration keeps the branch MERGEABLE + low-conflict; it does NOT by itself prove the branch's own work is correct (each change still crosses §11.4.108 runtime-signature + §11.4.40 retest), and it is NOT the approval-gated back-merge TO trunk (that stays §11.4.167(I)/§11.4.40/§11.4.41 no-merge-until-approved). Classification: universal (§11.4.17) — the consuming project supplies its canonical-trunk ref name, cadence unit, post-merge smoke-gate command, and per-track worktree/CoW mechanism per §11.4.35. Composes §9/§9.2/§11.4.6/§11.4.17/§11.4.37/§11.4.41/§11.4.42/§11.4.71/§11.4.84/§11.4.88/§11.4.89/§11.4.103/§11.4.113/§11.4.167/§11.4.176/§11.4.178/§11.4.179/§11.4.181. Propagation gate `CM-COVENANT-114-188-PROPAGATION` (literal `11.4.188`) + recommended gate `CM-REGULAR-MAIN-MERGE-INTO-FEATURE` (every live feature branch / track has integrated the canonical trunk within the declared cadence; its last integration is a real `git merge` not a rebase; the post-merge smoke + zero-conflict-marker + no-lost-commit checks passed) + paired §1.1 mutation (let a live branch go stale past the cadence, OR record a rebase of a shared branch in place of a merge, OR commit a conflict marker → the gate FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.188. Non-compliance is a release blocker. No escape hatch — no `--skip-main-merge`, `--merge-only-at-end`, `--rebase-shared-branch`, `--defer-trunk-sync`, `--no-post-merge-smoke`, `--force-push-after-merge`, `--big-items-only-trunk-sync` flag.

§11.4.189 — Most-reopened cases get extra-depth live-testing scrutiny FIRST (User mandate, 2026-07-10). Verbatim operator mandate: "All live testing MUST pay extra attention to cases which have been reopened numerous times and those cases in depth retested, investigated, fully validated and verified with real physical results and no bluff of any kind anywhere!" All LIVE TESTING (on-device / on-target runtime validation, deploy-and-observe) MUST give EXTRA-DEPTH retest + in-depth investigation + FULL validation/verification with REAL PHYSICAL captured evidence (§11.4.5/§11.4.69/§11.4.107 — captured audio/video/sysfs/dumpsys/sink-side/runtime-signature) and NO bluff of any kind ANYWHERE to the cases that have been REOPENED THE MOST TIMES (highest §11.4.55 reopens-count) — the empirically-most-fragile set gets the DEEPEST live scrutiny, and it gets it FIRST.

(1) **WHAT** — for the highest-reopens-count cases, live testing is NOT a single pass: it re-runs the case at EXTRA DEPTH (more iterations per §11.4.50 deterministic consistency, wider edge-case + regression coverage), RE-INVESTIGATES the underlying defect per §11.4.102 systematic-debugging + §11.4.146 reproduce-first, and CONFIRMS with rock-solid captured physical evidence per §11.4.123 — never metadata-only / config-only / absence-of-error / grep-without-runtime (§11.4/ §11.4.1), never a false result. (2) **ORDER** — the most-reopened set is scrutinised FIRST, ahead of the rest of the live-test suite (the §11.4.132 risk-descending discipline, with most-reopened the load-bearing key). (3) **KEY** — priority is keyed off the §11.4.55 reopens-count (the consuming project supplies its reopens-count source per §11.4.35, e.g. the §11.4.93 workable-items DB `reopens_count`); a high reopens-count is the strongest empirical fragility signal (§11.4.132(d)). (4) **PROOF** — each most-reopened case's live PASS cites its captured-evidence artefact path (§11.4.69/§11.4.107) on a CLEAN/fresh deployment (§11.4.108 runtime-signature), with the §11.4.115 RED → GREEN polarity flip captured where a permanent regression guard exists (§11.4.135). STRENGTHENS/REFINES §11.4.132 (risk-ordered validation factor (d) most-reopened — §11.4.189 promotes most-reopened into its OWN EXTRA-DEPTH mandate at the LIVE-TEST layer, not merely one ranking factor among four) + §11.4.55 (reopens-count as the priority key) + §11.4.129/§11.4.130 (validate-the-just-fixed / highest-risk-first after redeploy). Classification: universal (§11.4.17) — references no project-specific hardware; the consuming project supplies its reopens-count source per §11.4.35. Composes §11.4.5/§11.4.6/§11.4.55/§11.4.69/§11.4.107/§11.4.108/§11.4.115/ §11.4.129/§11.4.130/§11.4.132/§11.4.135/§11.4.146. Propagation gate `CM-COVENANT-114-189-PROPAGATION` (literal `11.4.189`) + recommended gate `CM-MOST-REOPENED-EXTRA-DEPTH-LIVE-TEST` (the live-test suite orders + extra-depth-scrutinises the highest-reopens-count cases FIRST, each with captured physical evidence on a clean deployment) + paired §1.1 mutation (drop the most-reopened extra-depth pass, OR run it AFTER the rest of the suite, OR PASS a most-reopened case on metadata-only evidence → the gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.189. Non-compliance is a release blocker. No escape hatch — no `--skip-most-reopened-depth` , `--reopened-same-as-any` , `--live-test-any-order` , `--metadata-pass-for-reopened-OK` , `--defer-reopened-scrutiny` flag.

§11.4.190 — Website engineering-quality mandate: every project website MUST be fully responsive + completely SEO-optimized + uniquely OpenDesign-authored + bleeding-edge enterprise-quality, each PROVEN with captured evidence (User mandate, 2026-07-10). Verbatim operator mandate: "every website we create inside our projects MUST BE fully responsive - working on all browsers, all operating systems, all platforms and all devices and screen sizes! Every Website MUST BE completely SEO optimized! Layouts and templates done MUST BE unique and everything fully designed from OpenDesign! Make sure this applies to our project Website as well! Everything must be bleeding-edge enterprise quality, stunning visual look and feel, innovative, creative and unique! There MUST BE no false or faulty result or bluff of any kind anywhere!" Every website / web-UI surface a project ships — INCLUDING this project's OWN website — MUST satisfy AND PROVE with captured physical evidence (§11.4.5/§11.4.69/ §11.4.107/§11.4.170 — never a bare claim, §11.4.6) ALL of: **(A) FULL RESPONSIVENESS** — correct, usable rendering across ALL major browser engines (Chromium/Firefox/WebKit), all operating systems, all device classes (desktop/laptop/tablet/phone/large-display) and the full range of screen sizes/ orientations/pixel-densities; NO broken layout / overflow / clipping / overlap / off-screen or unusable control at ANY breakpoint. **(B) COMPLETE SEO OPTIMIZATION** — semantic HTML + correct document outline; per-page unique `<title>` + meta description; Open Graph + Twitter cards; canonical URLs; structured data (schema.org / JSON-LD); `robots` + XML sitemap; crawlable / indexable markup; WCAG AA accessibility + Core Web Vitals performance (both feed ranking). **(C) UNIQUE OpenDesign-AUTHORED templates/layouts** — EVERY layout + template designed FROM the OpenDesign design-token system (§11.4.162), NOT generic / off-the-shelf / templated; visually unique, innovative, creative, brand-distinct. **(D) BLEEDING-EDGE ENTERPRISE VISUAL QUALITY** — stunning, polished, professional look-and-feel in BOTH light AND dark themes (§11.4.162 tokens). **(E) ANTI-BLUFF PROOF** — NONE of (A)–(D) may be CLAIMED without captured physical evidence: responsiveness PROVEN by device-independent host-rendered screenshots across the real breakpoint × engine matrix + an OCR/vision layout oracle (no overlap / label-over-label / clip / overflow / off-screen); SEO PROVEN by an automated audit (Lighthouse SEO + structured-data validation + meta / OG / sitemap / robots presence) meeting a defined score floor; visual quality PROVEN by the §11.4.170 device-independent host-rendered pixel proof per screen × state × {light,dark}; uniqueness PROVEN by OpenDesign-token provenance. A website reported responsive / SEO-optimized / unique /

enterprise-quality WITHOUT these captured proofs is a §11.4 PASS-bluff at the web-UI layer. §11.4.190 is the strict WEB-quality expansion of §11.4.162 (OpenDesign — it applies §11.4.162's design-token mandate to the full responsive + SEO + enterprise-quality web surface) and binds §11.4.170's device-independent host-rendered pixel proof as the correctness oracle; sibling of §11.4.168 (exported-document independent visual validation) at the web-UI layer. Classification: universal (§11.4.17) — references no project-specific hardware; the consuming project supplies its concrete breakpoint × engine matrix, SEO score-floor, and OpenDesign token set per §11.4.35. Composes §11.4.5/§11.4.6/§11.4.69/§11.4.107/§11.4.162/§11.4.168/§11.4.170. Propagation gate `CM-COVENANT-114-190-PROPAGATION` (literal `11.4.190`) + recommended gates `CM-WEBSITE-RESPONSIVE-PROVEN` (host-rendered screenshots across the breakpoint × engine matrix + layout oracle — no overlap/clip/overflow/off-screen) / `CM-WEBSITE-SEO-OPTIMIZED` (automated SEO audit meets the score floor + structured-data / meta / OG / sitemap / robots present) / `CM-WEBSITE-OPENDSIGN-UNIQUE` (every layout/template's OpenDesign-token provenance, no generic/off-the-shelf template) + paired §1.1 mutation (report a website responsive/SEO/unique/enterprise-quality with NO captured proof, OR ship a generic non-OpenDesign template, OR skip the light/dark host-rendered pixel proof → the recommended gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.190. Non-compliance is a release blocker. No escape hatch — no `--skip-responsive-proof`, `--seo-optional`, `--generic-template-OK`, `--visual-quality-later`, `--claim-without-render` flag.

§11.4.191 — Work-to-track/branch binding enforcement: a logic-group's work can NEVER be committed OR dispatched onto the wrong track/branch (research-derived, 2026-07-10). Compact summary: every logical group of workable items binds to EXACTLY ONE canonical (`track`, `branch`) recorded ONCE in the single source of truth (the §11.4.93 workable-items DB `logic_groups.destination` = authoritative branch + a `canonical_track` column + a `group_paths` file-scope manifest), and NO change belonging to a group may LAND on, nor be DISPATCHED to, a track/branch other than that group's canonical one — the file-scope generalisation of §11.4.181/§11.4.182 from branch-NAME consistency to work-PLACEMENT consistency (forensic FACT 2026-07-10: `srcpy/Remote-app` work — group `mistiq-vader`, canonical `feature/mistiq-vader` /T2 — was committed onto `main` /T1; §11.4.181's

guard-branch-consistency.sh never fired because committing onto an existing branch is not a feature-branch CREATE, and no gate checked which FILES land on which BRANCH). (1) **Single source of truth** — the logic_group → (branch, track, path-globs) binding is recorded ONCE in the DB and LOOKED UP, never re-invented (a commit/dispatch placing a group's file-scope on a non-canonical branch/track is a §11.4.6 no-guessing violation). (2) **Preventive hook (§11.4.109/§11.4.177)** — a PreToolUse guard guard-work-track-binding.sh (inherited BY REFERENCE, never copied) BLOCKS (exit 2) a git-commit-class Bash call whose to-be-committed file set contains a path owned by a group whose canonical branch/track ≠ the current checkout, AND BLOCKS an Agent/Task dispatch whose §11.4.182 label track/branch ≠ the canonical (track, branch) of the ticket's group; fail-closed on unreadable registry, # guardrails:allow <reason> audited escape. (3) **Detective gate (§11.4.110/§11.4.108)** — CM-WORK-TRACK-BINDING-ENFORCED asserts, on the current branch's own added commits (merge-base(HEAD,main)..HEAD), that every group_paths -owned file's group destination == current branch (and canonical_track == current track when pinned) — the actual enforcement boundary regardless of how the commit was made, the runtime signature being FAIL-on-wrong-placement / PASS-on-canonical-placement on a clean checkout. (4) **Risk-free reconciliation (§9.2/§11.4.113/§11.4.84)** — a mis-placed change is reconciled by MERGING its commits onto the group's canonical branch (union, no commit lost, no -s ours /reset), ff-only push to all upstreams (NEVER force-push), then removing it from the wrong branch only after the merge is confirmed on all mirrors. Classification: universal (§11.4.17) — references no project hardware; the consuming project supplies its (track, branch) map, group_paths seed, and commit-wrapper wiring per §11.4.35.

Composes

§11.4.6/.28/.29/.35/.58/.75/.84/.93/.95/.109/.110/.113/.119/.167/.176/.177/.178/.181/.182/ §9.2. Propagation gate CM-COVENANT-114-191-PROPAGATION (literal 11.4.191) + recommended gate CM-WORK-TRACK-BINDING-ENFORCED + paired §1.1 mutation.

Canonical authority: constitution submodule Constitution.md §11.4.191. Non-compliance is a release blocker. No escape hatch — no --allow-cross-track-commit , --skip-work-binding , --commit-any-branch-OK , --dispatch-any-track-OK , --no-file-scope-check flag.

§11.4.192 — Continuous multi-track auto-backfill: a FREE track MUST be IMMEDIATELY re-assigned its next-highest-priority domain work-set, never left idle while its domain has actionable items (User mandate, 2026-07-11).

Verbatim operator mandate: "once a Track is free, we shall assign next set of workable items to it ... full continuous multi-track development ... no stopping ... this is MANDATORY". Compact summary: the moment any track's worker COMPLETES its current work-set (its item(s) closed with anti-bluff evidence, or merged), that FREE track MUST be IMMEDIATELY assigned the NEXT-highest-priority actionable work-set from its OWN domain (its §11.4.191 logic_group / canonical branch) — a track is NEVER left idle while its domain still holds actionable workable items, and the conductor/ruler CONTINUOUSLY re-assigns with NO stopping (full continuous multi-track development); PROMOTES the §11.4.103(B) ≥ 3 -parallel-background-streams-with-auto-backfill invariant into the multi-track (§11.4.176/§11.4.178/§11.4.179/§11.4.187) layer (§11.4.103(B) backfills a STREAM the moment it is done; §11.4.192 backfills a TRACK the moment its worker is done, keyed to that track's canonical (track, branch) domain per §11.4.191). (1) **WHAT** — on any track-worker completion the conductor/ruler (§11.4.187) MUST, per that track's domain, look up the NEXT actionable work-set, CLAIM it exactly-once (§11.4.176 exactly-once claim registry, no two tracks claiming one item), and DISPATCH it as the track's new §11.4.182-labelled work stream — without waiting for a scheduled wake or an operator prompt. (2) **PRIORITY** — the next work-set is the track domain's highest-priority actionable item (§11.4.42/§11.4.72 priority order, audio-first where in scope), never an arbitrary pick. (3) **SOURCE-SIDE ADVANCES EVEN WHEN VALIDATION IS GATED** — when the next item's on-device VALIDATION is blocked on a pending build / device availability (§11.4.108 clean-target / §11.4.185 manual-QA / operator-gated flash), the track MUST still advance the item's SOURCE-SIDE work (investigation, RED-test §11.4.115, source patch, pre-build gate + paired §1.1 mutation, docs) and honestly §11.4.3/§11.4.52 SKIP-with-reason ONLY the genuinely device-gated validation portion — a build/device gate on the validation step is NOT a reason to idle the whole track. (4) **HONEST BOUNDARY (§11.4.6)** — a track idles ONLY when its domain GENUINELY has zero remaining actionable workable items OR every remaining item is externally blocked with no source-side advancement possible (§11.4.94(A)/§11.4.97/§11.4.101 idle-only-when-blocked closed set); "the merges are done" / "the current item is done" / "nothing visible right now" is NEVER a valid idle reason while the domain still holds progressable items (a §11.4.94+§11.4.97 violation); an unavoidable per-item block PARKS one item, it does NOT stop the track (§11.4.101 — the track backfills its next non-blocked domain item). Classification: universal (§11.4.17) — references no project hardware; the consuming project supplies its track → domain map + priority source

per §11.4.35. Composes

§11.4.42/.58/.72/.94/.97/.101/.103/.167/.176/.178/.179/.182/.187. Propagation gate

`CM-COVENANT-114-192-PROPAGATION` (literal `11.4.192`) + recommended gate `CM-CONTINUOUS-MULTITRACK-BACKFILL` (on a track-worker completion the ruler re-assigns that track's next-highest-priority domain work-set; a track with a completed worker AND remaining actionable domain items but no new dispatch → the gate FAILs) + paired §1.1 mutation (let a free track sit idle while its domain holds an actionable item, OR treat "merges done" as a valid idle reason → the gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.192. Non-compliance is a release blocker. No escape hatch — no `--allow-idle-track`, `--skip-backfill`, `--stop-multitrack`, `--merges-done-is-idle`, `--defer-track-reassign` flag.

Appendix C — Multi-upstream push

This constitution submodule is hosted on multiple Git providers, and consuming projects often are too. The submodule ships an `install_upstreams.sh` helper (or invokes the `install_upstreams` system command from the parent toolkit) that reads `Upstreams/*.sh` declarations and configures all remotes locally.

Every commit to the constitution submodule MUST be pushed to ALL configured upstreams. Use a shell helper that iterates `git remote` and pushes one by one, OR configure `origin` with multiple push URLs (`git remote set-url --add --push origin <url>` per remote).

§11.4.193 — Anti-blind-typing mandate: every UI interaction MUST be "seen" and "understood" via OCR/vision/screenshot proof, NEVER blind-typed (User mandate, 2026-07-13). Verbatim operator mandate: "No blind typing of any kind anywhere!!!" and "Login to apps and services MUST BE done with understanding and seeing the whole UI ALWAYS!!!" and "Add into the constitution that blind typing

is STRICTLY FORBIDDEN, and instead everything MUST BE seen and understood in coordination with proper models and means we have incorporated in HelixQA!!!" and "No bluff of any kind anywhere!!!"

Blind typing — sending keystrokes, text, or input events to a UI surface WITHOUT first capturing AND understanding what is on screen (OCR / vision / uiautomator / accessibility-tree / screenshot) AND confirming the result after each interaction — is STRICTLY FORBIDDEN everywhere and always, for every agent, every tool, every automation, every platform, every context. Blind typing is LITERAL bluff: the agent claims "credentials entered" or "login completed" without visual proof the keystrokes landed on the intended fields, on the intended screen, producing the intended result. A credential typed into the wrong field, onto the wrong screen, or into a WebView that swallowed the input silently is a §11.4 PASS-bluff at the UI-interaction layer — the agent reported success while the end-user feature (logged-in app) is broken.

(1) SEE-BEFORE-YOU-TYPE. Before ANY input text, input keyevent, adb shell input, uiautomator gesture, or equivalent UI-driving command, the agent MUST capture the current screen state via screencap + OCR OR uiautomator dump OR accessibility-tree dump OR vision-model frame AND confirm the intended target element is visible AND focused. The captured screen state IS the evidence per §11.4.5 and §11.4.69 feature class touch_input.

(2) VERIFY-AFTER-YOU-TYPE. After EVERY interaction step, the agent MUST capture a NEW screen state and confirm the intended change occurred. A sequence of 5 input text commands with a single screenshot at the end is blind typing — each step needs its own verification.

(3) OCR/VISION ORACLE MANDATORY. The verification MUST use at least ONE of: Tesseract OCR on a screencap at 150+ DPI, uiautomator dump with grep for expected text content, HelixQA VisionEngine frame analysis per §11.4.159 and §11.4.160, or accessibility-tree dumpsys confirming element state.

(4) WEBVIEW SURFACES. When uiautomator returns near-empty hierarchy (common in streaming-app login WebViews), the agent MUST fall back to OCR via screencap then tesseract per §11.4.117. If even OCR is infeasible (secure surface captures black, DRM-blanked), the login is genuinely operator_attended only — SKIP-with-reason per §11.4.3, NEVER blind-type-and-claim-success.

(5) CREDENTIAL-SAFE OCR. When OCR output may contain credentials, the agent MUST redact credential characters from OCR output and log files per §11.4.10. Evidence paths cited via `ab_pass_with_evidence`; evidence file is the redacted screenshot or OCR text with credential characters replaced.

(6) ANTI-BLUFF. Every login/interaction PASS MUST cite a captured-evidence artefact path containing pre-interaction screen state, post-interaction screen state, and OCR/vision verification showing the expected change. A PASS without these artefacts is a §11.4 PASS-bluff regardless of exit code. Self-validated per §11.4.107(10): the OCR/vision oracle MUST ship a golden-good and golden-bad fixture pair.

Classification: universal (§11.4.17). Composes §11.4 / §11.4.5 / §11.4.6 / §11.4.69 / §11.4.107 / §11.4.117 / §11.4.159 / §11.4.160 / §11.4.10 / §11.4.52 / §11.4.3.

Propagation gate CM-COVENANT-114-193-PROPAGATION (literal 11.4.193) + recommended gate CM-NO-BLIND-TYPING + paired §1.1 mutation (strip OCR verification before an input text → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md §11.4.193.

Non-compliance is a release blocker regardless of context. No escape hatch — no ¹¹blind-typing-ok, ⁴skip-ocr-verification, ¹⁹⁴assume-input-landed, login-without-proof, type-without-seeing flag exists.

2026 07 14

§ . . — Exhaustive all-scenario, all-angle code-review + verify-against-captured-runtime-evidence mandate (User mandate, — —)

Forensic anchor — verbatim user mandate (2026-07-14):

"CRITICAL: Code review MUST TAKE into the accounts all scenarios, look at the problem or change from all angles, deepest possible analysis! No gaps, shortcomings or false results or bluff is allowed! Extend our root constitution Submodule and all relevant files properly so the review process is more precise and we do not have situations like the last one!"

Forensic case study (genericised, 2026-07-14 — "the last one"). A code-review (§11.4.125) returned a clean **GO with ZERO findings** on a boot-crash fix. It dismissed a defensive clamp branch as an INFORMATIONAL NIT — "effectively unreachable ... the input is always valid" — WITHOUT PROVING that assumption against the actual runtime input space. The fix shipped, was deployed to the target, and the target STILL crash-looped: the degenerate input the review had assumed could not happen DID happen at runtime (an empty / default-initialised profile opened with the exact zero / default value the review assumed impossible). The crash was a MULTIPLICATIVE size computation — a product $A \times B$ (e.g. `frame_size = channels × bytes_per_sample(format)`) — with TWO INDEPENDENT zero-triggers; the review verified only ONE factor (`A` , e.g. `channels`) and ASSUMED THE OTHER (`B` , e.g. `format`) away. The crash-loop reproduced on the deployed target under conditions PRESENT in the pre-fix crash evidence that the review never cross-checked the fix against — the evidence was available and not consulted.

Every code-review this Constitution governs (STRENGTHENING §11.4.125 / §11.4.134 / §11.4.142 — the review remains NECESSARY; §11.4.194 makes it EXHAUSTIVE-and-evidence-grounded so a clean GO can no longer be returned on an unproven assumption) MUST, before returning GO:

(1) Enumerate the FULL input / scenario space of BOTH the change AND the defect it addresses — every dimension, every code path, every interacting feature, every input domain the change touches — NOT merely the reported / happy-path scenario. For a MULTI-FACTOR failure mode (e.g. a product of `N` terms, ANY of which being degenerate — zero, empty, null, default-initialised, out-of-range — triggers the failure), the review MUST enumerate AND INDEPENDENTLY VERIFY the fix covers EACH factor / dimension. Verifying ONE factor and ASSUMING the rest are safe is the EXACT forbidden gap this anchor closes (the forensic `channels` -verified / `format` -assumed pattern).

(2) Prove every assumption. "This can't happen" / "unreachable" / "always valid" / "the caller never sends X" / "the input is always in range" MUST be PROVEN from the code + spec + captured evidence — NEVER accepted on faith, NEVER inferred from "it usually is" (§11.4.6 no-guessing). An unproven assumption that later proves false is a §11.4 REVIEW-LAYER bluff of the SAME severity class as a PASS-bluff. A review NIT that DISMISSES a path / branch as unreachable / dead / defensive-

only MUST carry the PROOF of unreachability (the callers enumerated, the input domain bounded, the degenerate value shown genuinely impossible) — a bare "effectively unreachable" NIT with no proof is ITSELF a finding, not a dismissal.

(3) Verify the fix against ANY captured runtime evidence. When captured runtime evidence of the failure EXISTS — a crash log / tombstone / stack trace / core dump / sink-probe report / the §11.4.108 runtime-signature / a reproducer — the review MUST confirm the fix ACTUALLY PREVENTS the CAPTURED failure (walk the fix against the EXACT captured conditions, not a hypothetical). A review that GO's a fix which then FAILS at runtime on a scenario PRESENT in the evidence (or a scenario the review CHOSE NOT TO CONSIDER) is a bluff — the evidence was available and not cross-checked.

(4) All angles. The review MUST cover ALL of: correctness against the stated goal; the FULL input domain (every factor of a multi-factor computation, every boundary, every degenerate value); every downstream / interacting feature — ESPECIALLY must-not-break constraints (per §11.4.133 target / hardware safety, and any feature the change shares state / timing / a resource with — e.g. audio / multichannel / passthrough integrity); the TESTS' OWN coverage of every dimension (a test that checks only ONE factor of a multi-factor bug is ITSELF a coverage gap the review MUST catch and require closed); and the §1.1 paired mutation covering EACH dimension (a mutation that flips only one factor leaves the others un-guarded).

(5) Honest boundary (§11.4.6). "Exhaustive" means the ENUMERATED + EVIDENCE-GROUNDED scenario space — NOT a claim of literal omniscience. A dimension that genuinely cannot be analysed (external state the review cannot inspect, a runtime path with no reproducer) MUST be stated as an EXPLICIT gap (UNCONFIRMED: / UNKNOWN: / PENDING_FORENSICS: with a tracked follow-up) — NEVER silently assumed safe. A silently-assumed-safe un-analysed dimension is the §11.4.194 violation.

(6) The quality-harvest review classes (research-derived extension, 2026-07-23 — operator mandate: every check, diagnostic method, and preventive measure from the reopen/first-touch root-cause forensics is checked through EVERY code-review process). In addition to (1)–(5), every review MUST explicitly CHECK, per reviewed change: **(a) closure-evidence class vs defect layer (§11.4.226)** — no terminal or ready-class status in the change's scope moves on evidence whose class ranks below the defect layer's floor; a user-visible property "verified" by file-existence checks plus greps, an echo-assertion, or

a state-only repair closing a defect as fixed is a FINDING; **(b) needed measurements** — every load-bearing measurement the change or its tests performs satisfies §11.4.201(6)–(12) (a reported absence cites its class-matched control needle), and NO count was converted into a finding without reading the underlying lines — a COUNT IS A LEAD, THE LINES ARE THE FINDING (forensic FACT: twice in one measured cycle a count was read as state — a "mixed half-written doc" that was actually coherent with a labelled historical citation, and an inflated FAIL total that included the summary line counting itself); **(c) governance-and-gate integrity (§11.4.227)** — a change that NAMES a gate lands its implementation or a registered deferral in the same change, and a governance edit preserves anchor-block integrity (exactly-once per mirror, lockstep-identical, no anchor-number collision); **(d) reviewer-authored adversarial mutations** — the author's self-authored mutation set is NECESSARY but NOT SUFFICIENT validation of a guard/test: forensic FACT (from the consuming project's dated forensic corpus cited in the 2026-07-22 root-cause analysis): an author's six self-authored mutations were ALL caught while an independent reviewer's FIRST THREE different mutations ALL sailed through, each restoring a user-visible defect. The reviewer MUST attempt at least one mutation the author did not write (or record honest infeasibility per §11.4.3); a guard whose reviewer-authored mutation survives is unvalidated instrumentation per §11.4.115(F) and a FINDING.

Classification: universal (§11.4.17) — a platform-neutral review-rigor discipline reusable by ANY project; the consuming project supplies its concrete input registries, must-not-break constraint set, and captured-runtime-evidence sources per §11.4.35. Composes §11.4.125 (the code-review-agent gate this strengthens) / §11.4.134 (iterate-until-GO — a §11.4.194 gap is a finding that re-arms the loop) / §11.4.142 (every-change-reviewed — §11.4.194 raises the review's PRECISION bar) / §11.4.6 (no-guessing — the operative discipline behind clauses 2 + 5) / §11.4.108 (runtime-signature — the captured evidence clause 3 cross-checks against) / §11.4.130 (validate-the-fix-first-after-redeploy — the runtime confirmation that would have caught "the last one") / §11.4.107(10) (self-validated analyzer — the multi-factor coverage discipline applied to the review's OWN test-coverage check). Propagation gate `CM-COVENANT-114-194-PROPAGATION` (literal `11.4.194` across the consumer fleet) + recommended gate `CM-EXHAUSTIVE-REVIEW-ALL-SCENARIOS` (a review GO must record the enumerated scenario / factor space + the proof of every dismissed-as-unreachable branch + the cross-check against captured runtime evidence where it exists) + paired §1.1 meta-test mutation (strip

the enumeration / the assumption-proof / the captured-evidence cross-check from a review record → the gate FAILs; strip the literal → the propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.194.

Non-compliance is a release blocker regardless of context. No escape hatch — no

```
--partial-scenario-review , --assume-unreachable-OK ,  
11 4 195  
--skip-captured-evidence-crosscheck ,  
--single-factor-review-suffices , --dismiss-branch-without-proof flag.  
2026 07 14
```

**§ . . — Branch-structure governance:
taxonomy + merge-after-live-QA + flavor/product
non-merge (User mandate, — —)**

Forensic anchor — verbatim operator rules (2026-07-14):

1. Branch taxonomy: `feat/<NAME_OF_FEATURE>` , `product/
<NAME_OF_PRODUCT>` , `flavor/<NAME_OF_FLAVOR>` (extends §11.4.181 / §11.4.167(D)).
2. Tracks on FEATURE branches or `main` MUST — after full confirmation + validation + verification + full LIVE MANUAL TESTING (§11.4.185) — MERGE their work to `main` COMPLETELY (all submodules + main repo), ff-only, NEVER force-push (§11.4.113), no commit lost (§9.2) (extends §11.4.167(I) / §11.4.188).
3. Tracks with an entirely-new FLAVOR/PRODUCT do ALL work on a flavor branch AND/OR product branch (both may exist) — this product-line work does NOT merge to `main` the same way.
4. Commit + push ALL branches to ALL upstreams (§2.1), ff-only.

Every controlled branch in a governed repository MUST belong to one of THREE canonical taxonomy classes, and each class carries a distinct integration lifecycle:

(A) Taxonomy (extends §11.4.181 / §11.4.167(D)). A branch's PREFIX declares its class and its lifecycle:

Prefix	Class	Integration lifecycle
feat/<NAME_OF_FEATURE>	FEATURE	merges to main after full live-manual-QA (clause B)
product/ <NAME_OF_PRODUCT>	PRODUCT line	does NOT merge to main the standard way (clause C)
flavor/<NAME_OF_FLAVOR>	FLAVOR line	does NOT merge to main the standard way (clause C)

The <NAME_*> slug is lowercase snake/kebab (§11.4.29). The canonical branch NAME for a feature / logical group is minted ONCE and used IDENTICALLY across the main repo AND every touched owned submodule (§11.4.181) — recorded once in the §11.4.176 exactly-once claim registry / the §11.4.93 workable-items logic_group → branch mapping and LOOKED UP, never re-invented (§11.4.6). An untouched submodule stays on its base branch. The prefix STRINGS above are the operator-canonical form; a project already using a different feature-prefix convention (e.g. feature/...) is compliant so long as that convention is the ONE canonical name for the group per §11.4.181 — reconciling an inconsistent set of names is the §11.4.181 risk-free MERGE-onto-canonical path (ff-only, no commit lost, NEVER a rename that drops commits), never a silent rename (§11.4.6).

(B) FEATURE / main tracks — MERGE-TO-MAIN-AFTER-LIVE-QA. A track working on a feat/* branch (or on main directly) MUST, before merging its work to main, satisfy IN ORDER: full confirmation + validation + verification (automated four-layer per §11.4.4(b) + §11.4.40 full-suite retest) THEN full **LIVE MANUAL TESTING** by the QA team (§11.4.185 — the FINAL human sufficiency gate; automated GREEN is necessary but NOT sufficient). Only after §11.4.185 confirmation may the merge proceed, and it MUST be COMPLETE — the main repo AND every touched owned submodule — integrated **ff-only** onto the latest main (§11.4.113 merge-onto-latest-main), **NEVER force-push** (§11.4.113), with **NO commit lost / NO history rewrite** (§9.2 / §11.4.41 step 3: union, zero conflict markers, zero dropped files), pushed to **ALL upstreams** (§2.1, ff-only). This is §11.4.167(I) (no-merge-until-approved) + §11.4.188 (regular trunk-integration keeps the branch mergeable) — §11.4.195 binds the §11.4.185 live-manual-QA gate as an explicit precondition of the approval-gated back-merge.

(C) FLAVOR / PRODUCT lines — DO-NOT-MERGE-TO-MAIN-THE-SAME-WAY.

A track producing an ENTIRELY-NEW flavor and/or product line (a distinct downstream product, not a feature destined for the trunk) does ALL its work on a `product/<NAME>` AND/OR `flavor/<NAME>` branch (BOTH may exist for one line — a product line may carry a product branch AND a flavor branch simultaneously). This product-line work does NOT flow back into `main` via the clause-B merge — it is a divergent line with its own release lifecycle. It STILL obeys §11.4.113 (no force-push), §9.2 (no commit lost), §2.1 (push all branches to all upstreams, ff-only), and §11.4.188 (it MAY still merge `main` INTO itself to stay close to trunk — trunk-into-line is permitted and encouraged; it is the line-into-trunk direction that clause B gates and clause C withholds for product/flavor lines).

(D) Push discipline (§2.1 / §11.4.113). Commit + push ALL branches (feature, product, flavor) to ALL configured upstreams, ff-only, NEVER force-push. New branch creation is inherently ff (a new ref loses nothing, §9.2). Branch creation from an existing HEAD (`git branch <name> <sha>`) is non-disruptive (does not touch the working tree / index / other refs) and MAY proceed even in a dirty working tree under §11.4.84 provided no commit is in-flight (`.git/index.lock` absent, no live `git-commit` / `pack-objects`).

Honest boundary (§11.4.6): this anchor governs branch STRUCTURE + integration DIRECTION; it does NOT by itself prove any branch's work correct (each change still crosses §11.4.108 runtime-signature + §11.4.40 retest + §11.4.185 manual-QA).

Classification: universal (§11.4.17) — a platform-neutral branch-structure discipline reusable across any multi-track project; it references no project-specific hardware / vendor / region / asset, and is a strict generalisation of §11.4.181 (one-feature-one-canonical-branch-name) + §11.4.167(D)/(I) (feature work-stream lifecycle) + §11.4.188 (regular main → feature merge cadence) + §11.4.185 (manual-QA final confirmation). Composes §2.1 / §9 / §9.2 / §11.4.6 / §11.4.17 / §11.4.29 / §11.4.35 / §11.4.40 / §11.4.41 / §11.4.84 / §11.4.93 / §11.4.113 / §11.4.167(D)/(I) / §11.4.176 / §11.4.178 / §11.4.181 / §11.4.185 / §11.4.188 / §11.4.191. Propagation gate `CM-COVENANT-114-195-PROPAGATION` (literal `11.4.195` across the consumer fleet) + recommended gate `CM-BRANCH-TAXONOMY-CLASS` (every controlled branch's prefix is one of `feat/` | `product/` | `flavor/` | the reserved bases `main` / `master` ; a `product/` | `flavor/` branch is NOT the source of a merge onto `main` ; a `feat/` | `main` back-merge cites a §11.4.185 manual-QA confirmation) + paired

§1.1 mutation (a `product/*` branch merged onto `main` , OR a `feat/*` back-merge with no §11.4.185 citation → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.195.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--any-branch-prefix-OK` , `--merge-product-to-main` , `--skip-live-qa-before-merge` , `--force-push-branch` , `--partial-submodule-merge` flag.

§ . . . 2026 07 — **Native-alias-first priority + per-alias real-signal limit/subscription tracking + auto-rebind-on-recovery + resource-detection-by-real-identity-not-substring-match (User mandate, — —)**

Verbatim operator mandate: use ALL operational claude-NATIVE aliases FIRST — as long as ≥ 1 native alias is operational (NOT session-rate-limited, NOT weekly-limit-reached, NOT expired-subscription) — before ANY provider alias; track per alias the session rate-limit / weekly-limit / subscription expiry+renewal from REAL signals (never faked, §11.4.6) and record when each becomes operational AGAIN; as soon as a higher-priority native recovers, rebind tracks to it; and fix the host-budget guard footgun where a `pgrep -f substring-match false-REFUSED` every worker spawn. The multi-track alias-binding + rate-limit/subscription-tracking layer of the §11.4.187 ruler orchestration MUST implement ALL of: **(A) NATIVE-ALIAS-FIRST PRIORITY** — every OPERATIONAL native (subscription-based `claudeN`) alias is selected before ANY provider alias, UNCONDITIONALLY, as long as ≥ 1 native is operational; the guarantee is STRUCTURAL — a native/provider CLASS partition scanned native-class-completely-before-provider-class (the §11.4.111 resolve-by-CLASS-not-by-file-position discipline), so a provider can NEVER win while a non-cooled native exists, REGARDLESS of roster file order; within a class the operator-mandated order (native equal-capability, then providers `deepseek` → `xiaomi` → `opencode` → `kimi` → ...) is authoritative decision-DATA (§11.4.28) never re-invented (§11.4.6). **(B) PER-ALIAS REAL-SIGNAL LIMIT + SUBSCRIPTION TRACKING** — an "operational?" state per alias derived from REAL captured signals, NEVER faked/guessed (§11.4.6): the §11.4.187 captured 429 signature (`isApiErrorMessage` + `apiErrorStatus:429` / `api_error_status:429`) sub-classified into a reason-CLASS closed set

{session | weekly | subscription} from real message markers (reset/weekly-limit/session-limit); each limited alias records its CLASS + an operational-again EPOCH, and the three classes have DIFFERENT windows — session (429 self-heals at the reset, short), weekly-limit-reached (~until the weekly reset, long), subscription-expired (INDEFINITE — operational-again is UNKNOWN until a REAL renewal signal, parked with a far-future sentinel; §11.4.6 NEVER guesses a renewal date, an operator/config supplies the real one). A limited alias is NEVER auto-selected until its epoch passes (session/weekly) or an explicit renewal/recovery signal clears it (subscription). **(C) AUTO-USE-ON-RECOVERY** — the moment a higher-priority native recovers (its window elapsed, or a renewal marked it operational), the mechanism rebinds tracks UPWARD to it (a promote -class operation keyed off a comparable priority-rank: native class before provider, then config order), PRESERVING each track's worktree AND device leases (the leases key on the STABLE track id — the alias changes, not the track; §11.4.119 / §11.4.187). Idempotent — a track already on the best operational alias is left untouched. **(D) RESOURCE-DETECTION BY REAL IDENTITY, NOT SUBSTRING-MATCH (the RB-02 guard footgun fix)** — a host-safety "heavy work in flight" guard MUST identify a real resource (a heavy build, a live worker) by its ACTUAL process command line read from the OS (/proc/<pid>/cmdline on Linux), NEVER by a bare pgrep -f REGEX substring match that false-matches a process merely MENTIONING a token without BEING it — a pattern-CARRIER (a process whose args literally quote the whole alternation pattern; a claude /worker whose prompt embeds the token; a launcher/monitor/test), a same-tool ENGINE process, or the guard's own process/parent (the §12.12 self/sibling pgrep footgun — forensic FACT 2026-07-13: the guard false-REFUSED EVERY §11.4.187 spawn on a host with ZERO real builds because a live claude worker's prompt quoted the pgrep pattern). The guard re-reads each matched PID's real cmdline and excludes those carrier/engine/worker/self classes; an UNREADABLE/vanished cmdline is conservatively counted as the real resource (REFUSE — the safe, reversible default, §11.4.101 / §12.8), so a genuine build is still caught. **(E) ANTI-BLUFF (§11.4/§11.4.108)** — every piece carries four-layer coverage (pre-build gate + on-host test + paired §1.1 mutation + captured evidence per §11.4.69), each with a §11.4.115 RED-polarity reproduction of the defect on the pre-fix artifact + §11.4.50 determinism; **honest boundary (§11.4.6)** — this anchor mandates the MECHANISM + the REAL-signal derivation of "operational"; genuine live per-subscription quota-isolation proof still requires live tokens exercising real rate-limit boundaries (an operator-scheduled acceptance window), never conflated with a

mock/hermetic contract GREEN. Classification: universal (§11.4.17). Composes §11.4.187 (the ruler orchestration this is the alias-binding + rate-limit layer of) / §11.4.182 (track+branch+alias identity label) / §11.4.176 (multi-track exactly-once claim + device-lock) / §11.4.111 (resolve-by-stable-name/CLASS not by index/file-position) / §11.4.101 (autonomous-decision-over-blocking — the bounded-park-when-all-cooled + the conservative-refuse safe default) / §11.4.103 (continuous parallel-stream routine) / §11.4.6 (no-guessing — REAL signals, never faked) / §11.4.108 (runtime-signature) / §11.4.115 (RED-polarity) / §11.4.119 (single-resource-owner — leases key on the stable track id) / §12.12 (RLIMIT_NPROC / process-footgun awareness — the pgrep self/sibling match is the same footgun class) / §12.6 / §12.8 (host-safety budget + concurrency — the guard the RB-02 fix protects). Propagation gate `CM-COVENANT-114-196-PROPAGATION` (literal `11.4.196`) + recommended gate `CM-ALIAS-PRIORITY-LIMIT-TRACKING` (native-first CLASS partition present; per-alias reason-class + operational-again tracking; auto-rebind-on-recovery; the build-guard resolves by real cmdline not bare substring-match) + paired §1.1 mutation (swap the native/provider priority bases so a provider outranks a native → the recommended gate FAILs; flip the guard's strict-cmdline-filter to the pre-fix bare `pgrep` → its RB-02 test FAILs; strip the literal `11.4.196` from a carrier → the propagation gate FAILs; gate-code for the recommended gate = separate work item).

(F) CONFIGURED ≠ IN USE — the mechanism MUST BE WIRED TO THE ACTUAL WORKER SPAWNS, and its binding registry MUST REFLECT THE LIVE WORKER SET (extension, 2026-07-15). Forensic FACT (2026-07-15): the alias-binding orchestrator was fully installed and its configuration valid, yet its binding registry was `EMPTY ( (none bound) )` while THREE live workers ran OUTSIDE it on provider-backed aliases — so the recovery `promote` operation rebound ZERO tracks. The mechanism existed and managed NOTHING. Therefore: (1) EVERY worker / agent spawn path MUST bind THROUGH the orchestrator (the §11.4.187 launcher) — a worker started outside it is UNMANAGED (never rebound on recovery, never rate-limit-tracked, never counted against the host budget) and is a §11.4.196 violation; (2) the mechanism's READINESS is defined by its RUNTIME SIGNATURE (§11.4.108), NOT by its configuration — **the set of rows in the binding registry MUST EQUAL the set of live workers** (every live worker appears with its bound alias; every registry row corresponds to a live-or-honestly-terminal worker); an empty / partial registry alongside live workers is a FALSE-READY state; (3) claiming the mechanism

"ready" / "in use" on the strength of a valid config + an installed script, WITHOUT that registry ↔ live-worker equality captured as evidence, is a §11.4 / §11.4.6 bluff at the orchestration layer — "configured" is NOT "in use"; (4) the recommended gate `CM-ALIAS-PRIORITY-LIMIT-TRACKING` therefore ALSO asserts the registry ↔ live-worker equality, with a paired §1.1 mutation (leave a live worker unregistered → the gate FAILs). Composes §11.4.187 (the spawn / resume substrate this binds), §11.4.198 (the always-on default this makes REAL rather than nominal), §11.4.108 (runtime-signature-as-definition-of-done), §11.4.147 (a respawned agent re-binds through the orchestrator too).

Canonical authority: constitution submodule `Constitution.md` §11.4.196.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--provider-before-operational-native`, `--fake-limit-state`, `--guess-renewal-date`, `--no-rebind-on-recovery`,
`114197--substring-match-build-guard`, `--skip-real-signal-classification` flag.

§ . . — Research / kicked-off-work completion mandate: every research effort MUST be driven to full, wired, verified completion or explicitly closed — never left un-wired in the backlog (User mandate, — —)

Forensic anchor — verbatim operator mandate (2026-07-15):

"all research work we start MUST BE fully finished, completely confirmed, validated and verified! It MUST NOT happen for it to stay at the backlog uncompleted and un-wired! Loss of requirements, failure to incorporate them is FORBIDDEN!!! Everything MUST BE fully completed! No exceptions!"

Every research effort, kicked-off improvement, design doc, spike, prototype, proof-of-concept, or partially-landed feature that a project STARTS MUST be driven to FULL completion — OR explicitly, evidence-backed, CLOSED. There is no third "sitting un-wired in the backlog" state. "We researched it / designed it / started it" is NEVER a terminal state; a `docs/research/*` document, a design doc, a spike branch, or a half-wired feature that is neither COMPLETED-and-wired nor explicitly-

closed is a §11.4.197 violation of the same severity class as a §11.4 PASS-bluff at the requirements-integrity layer — the requirement was accepted, work began, and then it silently evaporated. Loss of requirements and failure to incorporate started work are FORBIDDEN.

COMPLETED means (ALL must hold): (1) **implemented** — the code / artefact exists; (2) **WIRED** — integrated into the system AND used by default, not merely present in the tree behind a dead flag / an un-called function / an unreferenced module / a never-selected path (mere-presence-without-wiring is the §11.4.124 dead-code gap AND the §11.4.108 layer-2/3 SOURCE → ARTIFACT → RUNTIME gap — a change present in source that never reaches the running artifact, or reaches it but is never exercised, is NOT completed); (3) **confirmed + validated + verified** with captured physical evidence per §11.4.5 / §11.4.69 / §11.4.107 on a clean deployment per §11.4.108 runtime-signature — never a metadata-only / config-only / absence-of-error / grep-without-runtime claim (§11.4 / §11.4.1); (4) tracked to a terminal Status per §11.4.15 / §11.4.33.

Explicitly CLOSED means: the effort is deliberately terminated with a documented, evidence-backed reason drawn from the existing closed vocabulary — §11.4.90 `obsolete` (superseded-by-design-change / superseded-by-later-mandate / duplicate-of / not-reproducible / feature-removed), §11.4.112 `structurally-impossible` (won't-fix, with cited impossibility proof), or §11.4.122 operator-confirmed drop — NEVER a silent abandonment. "No time" / "lower priority" is `operator-blocked` per §11.4.21 (with the self-resolution-exhaustion audit) or a still-open tracked item, NOT a close.

Mechanical binding to the tracker (§11.4.93): every research effort / design doc / kicked-off improvement MUST map to at least one tracked workable item (ATM-NNN per §11.4.54) whose completion is enforced by the loop. A `docs/research/*` document (or equivalent design / spike artefact) with NO tracked implementing item, OR a tracked item that has been stalled un-wired past the project's cadence, is a §11.4.197 violation. The autonomous loop (§11.4.87 / §11.4.94 / §11.4.97 / §11.4.103 / §11.4.126 / §11.4.192) MUST treat every such un-wired research effort as an actionable queue item — driving it to COMPLETED or explicitly CLOSED is exactly the "keep working until the queue is genuinely empty" contract, never optional backlog.

Honest boundary (§11.4.6): "fully completed" is bounded by the completion + closure definitions above (wired + verified, OR evidence-backed closed) — it does NOT mean every speculative idea must ship; it means every STARTED effort reaches a documented terminal state, and the terminal state is proven, never assumed.

Classification: universal (§11.4.17) — a platform-neutral requirements-integrity discipline reusable by ANY project; the consuming project supplies its research-doc tree path + tracker binding per §11.4.35. Composes §11.4.8 (deep-web-research — the research this mandate insists on FINISHING, not merely starting) / §11.4.87 / §11.4.94 / §11.4.97 / §11.4.103 / §11.4.126 / §11.4.192 (the autonomous-loop queue-drain that drives un-wired efforts to completion) / §11.4.118 (discovery-pressure — enumerated coverage of what was started) / §11.4.123 (rock-solid-proof — completion PROVEN, never claimed) / §11.4.124 (dead/unwired-code investigate-before-remove — the wiring half) / §11.4.150 (deep multi-angle web research per change) / §11.4.90 / §11.4.112 / §11.4.122 / §11.4.21 (the explicit-close vocabulary) / §11.4.93 / §11.4.54 (the tracker binding) / §11.4.108 (wired-and-verified on a clean target) / §11.4.15 / §11.4.33 (terminal Status). Propagation gate `CM-COVENANT-114-197-PROPAGATION` (literal `11.4.197`) + recommended gate `CM-RESEARCH-COMPLETION-TRACKED` (every `docs/research/*` effort / design doc maps to an open-or-closed tracked item; no research doc lacks an implementing item; no tracked item is stalled un-wired past the cadence) + paired §1.1 mutation (add a `docs/research/*` effort with no tracked implementing item, OR mark a research effort "done" while its feature is present-but-unwired → the recommended gate FAILs; strip the literal → the propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.197.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--research-may-stay-unwired`, `--skip-completion-tracking`, `--abandon-without-close`, `--present-but-unwired-is-done`, `--backlog-research-OK`, `--lose-requirement-OK` flag.

§₂₆ 07 15 . . — Default working mechanisms: multi-alias native-first orchestration AND heavy token-optimization are ALWAYS-ON defaults for single- and multi-track work (User mandate, — —)

Forensic anchor — operator mandate (2026-07-15): multiple claude-toolkit aliases (native-first) AND heavy token-optimization MUST be the DEFAULT mechanisms, always used with single- and multi-track development and with single or multiple working agents.

Two proven mechanisms are promoted to STANDING, always-on defaults — engaged automatically for EVERY working configuration (single-agent OR multiple-agent; single-track OR multi-track), never a per-request opt-in:

(A) Multi-alias native-first orchestration is the DEFAULT. The claude-toolkit multi-alias mechanism — every OPERATIONAL claude-NATIVE alias selected before ANY provider per §11.4.196 (native-alias-first priority + per-alias real-signal limit/subscription tracking + auto-rebind-on-recovery), driven by the §11.4.187 automatic multi-track ruler orchestration — MUST be the default execution substrate for ALL work, NOT only for the multi-track case. A SINGLE-track or SINGLE-agent session STILL binds its worker through the same native-first alias-priority mechanism (so it uses the best operational native alias + rebinds on recovery), and a MULTI-track / multiple-agent configuration uses it across every track (§11.4.176 exactly-once claim + §11.4.182 track+branch+alias identity label). Running a session on an arbitrary / hardcoded / provider-first alias when an operational native exists — or bypassing the alias-priority mechanism because "it is just one agent" — is a §11.4.198 violation.

(B) Heavy token-optimization is the DEFAULT. The §11.4.141 token-efficiency discipline (compact-in-CLAUDE.md/AGENTS.md + full-one-hop-away anchor layout, targeted reads over whole-file dumps, subagent context isolation per §11.4.20 / §11.4.70, evidence-hash references not evidence dumps) MUST be applied ALWAYS — on every read, every dispatch, every doc edit — for single- AND multi-track work, with single OR multiple agents. Token-heavy working patterns (re-reading whole large files when a targeted slice suffices, dumping full evidence into context, monolithic non-subagent execution of parallelisable work) when a token-efficient path exists are §11.4.198 violations.

Honest boundary (§11.4.6): "default mechanism" means engaged automatically without an operator request; it does NOT weaken any host-safety bound — the multi-alias / multi-agent fan-out stays under the §12.6 60% memory ceiling + the §12.12 thread-headroom check + the §11.4.58 parallel-agent cap (spawning contending / redundant agents is NOT "using the default", per §11.4.183 maximum-USEFUL-throughput). §11.4.198 BINDS the two mechanisms as always-on; it does NOT redefine them — their mechanics live in §11.4.196 / §11.4.187 / §11.4.141.

Classification: universal (§11.4.17). Composes §11.4.141 (token-efficiency — mechanism B) / §11.4.187 (automatic multi-track ruler orchestration — the substrate) / §11.4.196 (native-alias-first priority + limit tracking — mechanism A) / §11.4.182 (track+branch+alias identity) / §11.4.176 (multi-track exactly-once claim) / §11.4.183 (maximal multi-agent utilization — the bounded-throughput sibling) / §11.4.103 (continuous parallel-stream routine) / §11.4.58 (parallel-development PWU) / §11.4.20 / §11.4.70 (subagent-driven default) / §11.4.126 (default autonomous-loop) / §12.6 / §12.12 (host-safety bounds). Propagation gate `CM-COVENANT-114-198-PROPAGATION` (literal `11.4.198`) + recommended gate `CM-DEFAULT-MECHANISMS-ALWAYS-ON` (the native-first alias-priority mechanism + the §11.4.141 token-optimization discipline are wired as the default for single- AND multi-track / single- AND multi-agent configurations) + paired §1.1 mutation (make provider-first / bypass-alias-priority the single-agent default, OR strip the token-optimization default → the recommended gate FAILs; strip the literal → the propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.198.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--single-agent-skips-alias-priority`, `--provider-first-default`, `--token-optimization-optional`, `--heavy-tokens-OK`, `--not-default-mechanism` flag.

§ . 2026 07 15 — Exact-reproduction-sequence
mandate: when a working reproduction exists, the
investigation **MUST** use ITS exact sequence — a
deviating repro proves **NOTHING** (research-
derived, — —)

Forensic anchor (FACT, 2026-07-15). An investigation concluded that a defensive "cover" mechanism "never engages" and that the defect it guards was therefore unfixed. The conclusion was **WRONG**: the hand-rolled reproduction drove the target with a **DIFFERENT** input than the existing test did (it pressed a **HOME** key where the test sent a media-pause key), so the hand-rolled path never reached the **PRECONDITION** the mechanism keys off. Re-driven with the **TEST'S EXACT** sequence, the mechanism provably **ENGAGED**. A deviating reproduction had produced a confident, captured-looking, and entirely false negative.

Every investigation / debugging pass / validation that reproduces a defect or exercises a code path **MUST** satisfy **ALL** of:

(1) Use the existing reproduction's EXACT sequence. When a working reproduction **ALREADY EXISTS** — a test, harness, script, or recorded sequence that demonstrably reaches the defect or the code path under investigation — the investigation **MUST** drive the target with **THAT EXACT** sequence: same inputs, same order, same timing, same preconditions, same environment / topology (§11.4.3). A hand-rolled approximation is **NOT** a substitute.

(2) A deviating variant MUST first PROVE it reaches the precondition. If a hand-rolled / simplified / alternative sequence **IS** used, it **MUST FIRST** be proven — with captured evidence — to reach the **SAME** precondition (the entry-state the mechanism or defect keys off) before **ANY** conclusion is drawn from it.

(3) A repro that misses the precondition proves NOTHING. It is **NOT** evidence of absence (§11.4.6 — absence of the observation is not observation of absence). Concluding "the defect is present" / "the defect is absent" / "the mechanism never engages" from a sequence that never reached the precondition is a §11.4 / §11.4.1 bluff at the **INVESTIGATION** layer — a **FAIL-bluff** when it wrongly condemns working code (and drives a needless "fix" into a healthy path), a **PASS-bluff** when it wrongly clears broken code.

(4) Record the exact driving sequence as captured evidence. Every reproduction MUST record the precise sequence it drove (commands / input events / order / timing) so it is replayable + auditable; any deviation from an existing reproduction MUST be explicit, justified, and precondition-proven.

STRENGTHENS §11.4.102 (systematic-debugging Phase 2 — reproduce), §11.4.146 (reproduce-first), and §11.4.115 (RED baseline on the broken artifact): those mandate THAT you reproduce; §11.4.199 mandates HOW — with the sequence PROVEN to reach the defect, not one that merely resembles it. Classification: universal (§11.4.17) — the consuming project supplies its concrete input / driving mechanism per §11.4.35. Composes §11.4.1 / §11.4.3 / §11.4.5 / §11.4.6 / §11.4.7 / §11.4.102 / §11.4.107 / §11.4.115 / §11.4.123 / §11.4.135 / §11.4.146. Propagation gate CM-COVENANT-114-199-PROPAGATION (literal 11.4.199) + recommended gate CM-EXACT-REPRO-SEQUENCE (an investigation's conclusion cites the exact driving sequence used AND, where an existing reproduction exists, that the same sequence was used OR that the deviating sequence reached the precondition with captured proof) + paired §1.1 mutation (record a conclusion drawn from a deviating sequence with no precondition proof → the gate FAILs; strip the literal → the propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md §11.4.199.

Non-compliance is a release blocker regardless of context. No escape hatch — no --approximate-repro-OK , --skip-precondition-proof , --hand-rolled-repro-suffices , --absence-proves-absence , --deviating-sequence-conclusion-OK flag.

2026 07 15

§ . . — **Non-targetable deploy/flash tooling MUST isolate exactly ONE eligible target and VERIFY-AFTER-WRITE on the INTENDED target (research-derived, — —)**

Forensic anchor (FACT, 2026-07-15). A flashing tool with no device-selection argument AUTO-SELECTED a device from the bus. With a sibling board still attached, it silently wrote the freshly-built image to the WRONG board, printed a SUCCESS message, and the INTENDED board kept its OLD build. Every

downstream test then validated a target that had never received the artifact: the §11.4.108 ARTIFACT → RUNTIME-ON-CLEAN-TARGET layer was silently broken while every tool and gate reported green.

Whenever a deploy / flash / install / write tool CANNOT address a specific target (no serial / id / path selector, or one that the procedure does not pass), the procedure MUST satisfy ALL of:

(1) Targetability is a PRECONDITION. Before the write, the procedure MUST establish that the intended target is the one that will be written: EITHER the tool takes an explicit stable target selector (serial / UUID / stable id per §11.4.111 — never an enumeration ordinal) and it IS passed, OR the procedure GUARANTEES that EXACTLY ONE eligible target is present on the bus / host / network at the moment of the write (every sibling disconnected, powered down, or excluded), with the enumerated eligible-target set captured as evidence (size == 1).

(2) The tool's own success message is NOT proof. A tool's "ok" / exit 0 proves only that bytes reached the device the TOOL selected — NEVER that they reached the device the OPERATOR intended (§11.4.6: "the write succeeded" ≠ "the write succeeded ON THE INTENDED TARGET"). Treating a tool's success line as deployment evidence is a §11.4 bluff at the deployment layer.

(3) VERIFY-AFTER-WRITE on the intended target. After the write, the procedure MUST READ BACK the deployed artifact's IDENTITY from the INTENDED target ITSELF (build id / version / checksum / artifact fingerprint) and assert it EQUALS the artifact just written. A mismatch — or the OLD identity still present — is a FAIL, never a warning. This is the §11.4.108 ARTIFACT → RUNTIME assertion applied to the deployment step, and it is the ONLY thing that closes the wrong-device write.

(4) No validation on an unverified target. Any test / validation run against a target whose deployed-artifact identity has NOT been verified is INVALID, and any PASS it produces is a §11.4 PASS-bluff — it exercised code that may never have been deployed there (the §11.4.108 clause-4 stale/wrong-artifact class).

Honest boundary (§11.4.6): isolation + read-back prove that the INTENDED target holds the INTENDED artifact; they do NOT prove the artifact is CORRECT (that remains §11.4.108 / §11.4.40 / §11.4.130). Classification: universal (§11.4.17) — the consuming project supplies its concrete deploy tool, target-identity read-back, and eligible-target enumeration per §11.4.35. Composes §11.4.6 / §11.4.46 / §11.4.108 / §11.4.111 (stable identifier, never an ordinal) / §11.4.119 (single-

resource-owner) / §11.4.130 / §11.4.133 (writing an image to an unintended device is a TARGET/HARDWARE-safety event) / §11.4.139. Propagation gate `CM-COVENANT-114-200-PROPAGATION` (literal `11.4.200`) + recommended gate `CM-DEPLOY-TARGET-ISOLATED-AND-VERIFIED` (an explicit stable selector was passed OR the eligible-target set was captured with `size == 1`; the post-write identity read-back from the intended target is present and matches) + paired §1.1 mutation (deploy with two eligible targets on the bus and no explicit selector, OR strip the post-write identity read-back → the gate FAILs; strip the literal → the propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.200.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--auto-select-target-OK`, `--trust-tool-success-message`, `--skip-post-write-verify`, `--validate-unverified-target`, `--multiple-targets-on-bus-OK` flag.

2026 07 15

§ . . — Every guard/gate MUST assert the REAL condition: a false-positive refusal is a FAIL-bluff, a false-negative pass is a PASS-bluff (research-derived, — —)

Forensic anchor (FACT, 2026-07-15 — two independent instances in one cycle). (a) A host-budget guard used a bare `pgrep -f '<pattern>'` and REFUSED every worker spawn on a host with ZERO real builds, because a live worker's own command line merely QUOTED the pattern (the §11.4.196(D) / §12.12 carrier footgun). (b) A host-safety pre-flight refused a build on a "swap full" reading that did not reflect a real memory-pressure condition, and cleared only after the operator manually cycled swap. In BOTH cases a guard refused on a PROXY signal that was NOT the real condition — halting real work while reporting a condition that did not exist.

Every GUARD / GATE / PRE-FLIGHT CHECK (host-safety, resource-budget, readiness, lock, capability, availability) MUST assert the REAL condition it CLAIMS to assert, resolved from the AUTHORITATIVE source — the process's real `/proc/<pid>/cmdline`, the real cgroup / meminfo pressure metric, the real lock-holder liveness (`kill -0`), the real device identity — NEVER from a proxy signal that something which is NOT the condition can satisfy.

(1) A FALSE-POSITIVE REFUSAL IS A FAIL-BLUFF (§11.4.1). A guard that BLOCKS work when the condition is ABSENT is exactly as forbidden as a gate that PASSES while a defect is present: it halts real work, it teaches operators and agents to bypass guards (the worst possible second-order effect), and it hides the true state behind a confident false claim.

(2) A FALSE-NEGATIVE PASS IS A PASS-BLUFF (§11.4). A guard that does NOT fire when the condition IS present is the classic bluff — the guard exists precisely to catch that state.

(3) SELF-VALIDATED (§11.4.107(10)). Every guard MUST ship a golden-TRUE fixture (the real condition PRESENT → the guard FIRES) AND a golden-FALSE fixture that INCLUDES the known CARRIER / decoy case (the condition ABSENT but the proxy signal present → the guard MUST NOT fire), both wired into the meta-test. A guard that fires on its golden-FALSE fixture is ITSELF the defect, and shipping a guard without both fixtures is a §11.4 bluff at the guard layer.

(4) CONSERVATIVE-SAFE DEFAULT ON AN UNRESOLVABLE SIGNAL (§11.4.101 / §12). When the authoritative source genuinely cannot be read, the guard takes the SAFE, reversible default (refuse for a destructive / high-blast-radius action; the safe side for a host- or hardware-safety guard) and SAYS SO HONESTLY. A conservative refusal WITH an honest "could not resolve the condition" reason is NOT a false positive; a refusal that CLAIMS a condition it never verified IS one (§11.4.6).

(5) EVERY REFUSAL REPORTS ITS RESOLVED EVIDENCE. A guard refusal MUST print the resolved evidence that triggered it (the real pid + cmdline, the real metric + threshold, the real lock holder + liveness result) so the refusal is auditable and a false positive is diagnosable in one step rather than one session.

(6) SCOPE — EVERY LOAD-BEARING MEASUREMENT, NOT ONLY GUARDS (extension, research-derived, 2026-07-17). Clauses (1)-(5) bind guards / gates / pre-flights — decision seams that REFUSE or ALLOW — and their failure taxonomy is refusal-shaped. But the same defect kills an ad-hoc measurement that refuses nothing: a `grep`, a count, a parse, an exit-code check, a string scan, a burn-down metric, an artifact inspection. These are the instruments an agent forms BELIEFS from, and a false belief propagates into every downstream decision without ever passing a guard. Therefore clauses (6)-(8) bind **every measurement whose result is load-bearing** — i.e. any observation an agent, gate, report, or

metric relies on to assert a fact. Their taxonomy is MEASUREMENT-shaped: a **FALSE-NULL** (a zero / empty / no-match / "absent" read as evidence of absence when the instrument simply could not see) and a **FALSE-MATCH** (a hit on something that MENTIONS the thing rather than IS it) are both §11.4/§11.4.1 bluffs at the measurement layer — the null one is the more dangerous, because a broken instrument and a genuinely-clean artifact return the identical, confident, quiet zero.

(7) CARRIER-VS-THING + THE CONTROL NEEDLE. (a) MATCH STRUCTURE, NOT SUBSTRING. A measurement **MUST** distinguish the THING from a CARRIER that merely MENTIONS it. Forensic FACTs, all genericised, all real, all in ONE cycle of ONE project: a removal's completeness was measured by counting references to the removed component — and the count matched the AUDIT-TRAIL COMMENTS documenting the removal, so a completed removal read as a failed one (the §11.4.122 audit trail NAMES what it removed; naming is not using); a "does the artifact contain X" scan matched the governance text DOCUMENTING X; and this anchor's own siblings had already booked the same shape twice — §11.4.196(D)'s process-carrier (a process quoting the pattern is not a process being the thing) and §12.12's self-match (the scanner's own command line contains the pattern it scans for). **A defect already anchored twice and still recurring is the argument for the general rule, not against it.** The instance that decides it: a mutation-residue scan run over a governance repository flagged the very anchor that DEFINES the mutation markers, because that anchor quotes them in its rule text — the enforcement mechanism false-positived on its own law. **(b) A NULL IS NOT EVIDENCE UNTIL A CONTROL NEEDLE PROVES THE INSTRUMENT CAN SEE.** Before a zero / empty / no-match / "absent" result may be reported as a real absence, the SAME instrument, over the SAME path, against the SAME artifact, **MUST** be shown to return a non-null for a **control needle** — a string/property **KNOWN** to be present. Needle NOT found ⇒ the instrument is blind and the zero says NOTHING about the artifact (§11.4.6: report the blindness, never the absence). **The needle MUST be expressed through the SAME load-bearing query features as the query it certifies** — same dialect constructs (alternation / character class / quantifier / escape), same anchoring, same quoting, same encoding path. A bare-literal needle certifies ONLY the layers a bare literal crosses; it does NOT certify a query whose meaning depends on a feature the needle never used. **This bound is not theoretical — the (7)(c) dialect FACT refutes the naive form directly:** a query intending alternation but written with an escape that the dialect treats as a LITERAL pipe searches a string nobody wrote and returns zero;

a plain-literal needle sails through the identical tool + path, returns non-null, and would "certify" that zero — while the query is still blind. A needle carrying the SAME alternation (one arm known-present) dies with the query and exposes it. Hence the earned inference is bounded: **needle found $\Rightarrow$ the path can see the needle's QUERY CLASS $\Rightarrow$ a zero from a query of that same class is evidence**; a zero from a query using features the needle did not exercise is NOT yet evidence, and needs a needle of its own class. This generalises §11.4.115(F)'s validate-the-detector principle (a guard never observed FAILing on the genuinely-broken artifact is unvalidated instrumentation and mints no verdicts) from guards minting verdicts to every measurement, and it is PER-QUERY, against the artifact in hand — **golden fixtures cannot substitute** (§11.4.146(D3): "goldens are necessary, not sufficient"), because an instrument can pass every fixture on the authoring host and still be blind on the real path. **(c) THE PATH IS PART OF THE INSTRUMENT.** The instrument is not the tool, it is the tool PLUS every layer the query crosses, and the load-bearing failures live in the layers: a quoted pattern whose escapes are consumed by an intervening/wrapper shell, so the far side searches a LITERAL where alternation was intended and matches nothing; `\|` under an extended-regex dialect where it is a LITERAL pipe, not alternation, silently searching for a string nobody wrote; a pipeline (`check | head , check | xargs`) whose exit status is the LAST stage's, so a checker's verdict is discarded and its "clean ✓" is the pager's success, repeatable indefinitely without ever checking anything; a pattern anchored to line-start that misses an indented or colour-escaped line, reading a present verdict as no-verdict; a binary/encoding-sensitive extractor that splits a multi-byte character and reports a present string as absent. Each of these returns a CLEAN, CONFIDENT, WRONG answer — none crashes, so §11.4.1's script-bug clause (which covers failures that crash into a FAIL) never fires, and §11.4.67's parse-check (which is parse-time only) cannot see it. Hence (7)(b) is not optional diligence: a class-matched control needle is the only mechanism that catches these defects, because it is the only one that exercises the whole path with the query's own features. **Honest boundary (§11.4.6):** the needle proves the path can SEE its own query class — it does NOT prove the query is semantically the RIGHT question to ask (a correct-and-visible query for the wrong property still misleads); that remains (8)'s and §11.4.201(1)'s concern.

(8) A METRIC MUST BE VALIDATED AGAINST THE DEFINITION OF DONE. A progress / burn-down / coverage / ratchet metric is a load-bearing measurement, and a metric that is ANTI-CORRELATED with done is not a misread condition — it

is the WRONG condition, read correctly. Forensic FACT (genericised): a removal's burn-down counted references to the removed component and was tracked toward zero — but the CORRECT end-state ADDS references, because a retired gate must carry a mandated rationale naming what was removed and why, and the inverse "stays-removed" assertion must NAME the removed thing to assert its absence; the number therefore RISES as the work succeeds, the target is unreachable, and chasing it to zero would DELETE the very rationale and assertions the constitution mandates. Therefore: before a metric may gate, rank, or report work, it MUST be checked against the definition of done — construct the correct end-state and evaluate the metric there; **if the correct end-state does not move the metric to the target, the metric is INVALID and MUST NOT gate work** (a metric whose target the correct outcome cannot reach will, if enforced, drive work AWAY from correctness — the strongest possible form of §11.4.201's proxy-signal defect, because here the proxy is not merely satisfiable by a non-condition, it is REFUTED by the condition). This is the (7)(a) carrier defect relocated from a guard to a metric: the audit-trail rationale and the stays-removed assertion are CARRIERS, and the burn-down was a substring-counter that carriers satisfy.

Honest boundary + standing obligation (§11.4.6): this constitution already ships at least one gating metric of exactly this class — §11.4.135's adoption ratchet, a monotone-decrease burn-down that BLOCKS the release tag (§11.4.135's seam; NOT the build) — and its validation against the definition of done is OWED, not hereby claimed; naming it is the §11.4.118 discovery-pressure obligation. That validation **MUST be tracked** as a §11.4.197 work item by every consuming project that operates a ratchet (stated as the obligation it is — this anchor does NOT assert that any particular project has already filed it; asserting an unverified tracker state would be the §11.4.6 violation this very clause exists to forbid). **Interim precedence (§11.4.6 — stated, never left for a consumer to guess):** until that validation lands, §11.4.135(5)'s ratchet KEEPS GATING — clause (8) does NOT retroactively disarm it, because an unvalidated ratchet that blocks a release is a bounded, visible, operator-resolvable stop, whereas disarming the only adoption mechanism would silently re-open the unguarded-backlog hole §11.4.135 exists to close (the §11.4.101 reversible-safe choice). Clause (8) binds (a) every NEW metric from this anchor forward — none may gate until validated against the definition of done — and (b) the ratchet's own validation as owed work, NOT as a licence to ignore it. A consuming project MUST NOT cite clause (8) to switch its ratchet off.

GENERALISES §11.4.196(D) (resource-detection-by-real-identity — from the alias / build guard to EVERY guard) and §11.4.180 (holder-liveness PROVEN via `kill -0` before reaping — to every guard's condition). Classification: universal (§11.4.17) — the consuming project supplies its concrete guards, authoritative sources, and carrier/decoy fixtures per §11.4.35. Composes §11.4.1 / §11.4.6 / §11.4.101 / §11.4.107(10) / §11.4.109 / §11.4.111 / §11.4.180 / §11.4.196(D) / §12 / §12.6 / §12.8 / §12.12. Propagation gate `CM-COVENANT-114-201-PROPAGATION` (literal `11.4.201`) + recommended gate `CM-GUARD-ASSERTS-REAL-CONDITION` (every guard resolves its condition from the authoritative source, ships golden-TRUE + golden-FALSE-with-carrier fixtures wired into the meta-test, and prints its resolved evidence on refusal) + paired §1.1 mutation (replace a guard's authoritative resolution with a bare substring / proxy match → its golden-FALSE carrier fixture makes the gate FAIL; strip the literal → the propagation gate FAILs; gate-code = separate work item).

Gate contract for clauses (6)-(8) (extension, 2026-07-17). `CM-GUARD-ASSERTS-REAL-CONDITION` is EXTENDED beyond its guard-fixture invariants to cover the measurement layer, and a metric-specific sibling is added — because a clause landing with no gate + no paired mutation is itself the §1.1 bluff-gate pattern this anchor exists to forbid: **(i)** every load-bearing NULL/ZERO/absence result reported as a finding cites a control-needle result of the SAME query class, captured alongside it (a reported absence with no class-matched needle evidence → FAIL); **(ii)** the needle's query class is asserted to match the certified query's load-bearing features (a bare-literal needle certifying an alternation/anchored/encoded query → FAIL — the finding-1 defect this clause was corrected for); **(iii)** recommended gate `CM-METRIC-VALIDATED-AGAINST-DONE` — every metric that gates, ranks, or reports work carries a recorded definition-of-done evaluation showing the correct end-state reaches the target (a gating metric with no such record, or one whose recorded correct-end-state does NOT reach target, → FAIL). Paired §1.1 mutations: swap a class-matched needle for a bare literal on a dialect-dependent query → (ii) FAILs; strip the needle entirely from a reported absence → (i) FAILs; register a metric whose correct end-state moves it AWAY from target → (iii) FAILs. Gate-code = a separate work item; this is the contract, not a claim that the code has shipped (§11.4.6).

(9) FIELD-IDENTITY + UNITS + SEMANTICS — capacity is not occupancy (extension, research-derived, 2026-07-23). Forensic FACT (genericised, 2026-07-22): a load-bearing latency claim — "a standing ~1.4-second send queue"

— was minted by reading a socket diagnostic's WINDOW-CAPACITY field (`rcv_wnd` 181120 bytes, the ceiling the peer ADVERTISES) as if it were QUEUE OCCUPANCY; the real observed Send-Q was a few hundred bytes (in-flight-sized, `app_limited`), three orders of magnitude smaller. The misread propagated into THREE parallel agents' briefs and into operator reports before anyone re-read the raw field; two INDEPENDENT re-derivations from the raw samples then converged on the correction, and the false figure had to be publicly withdrawn from a changelog. Therefore: before ANY metric read from a summary / diagnostic line / dashboard cell may gate a decision, rank work, or propagate into a brief / report / tracker, its FIELD IDENTITY (which field, exactly, in the authoritative raw output), its UNITS, and its SEMANTIC CLASS — capacity/limit vs occupancy/backlog vs rate vs cumulative counter vs high-water mark — MUST be verified against the authoritative source's own field definition (§11.4.99 latest-source where the tool's documentation defines it). CAPACITY≠OCCUPANCY is the named canonical instance; limit-vs-current, cumulative-vs-delta, and advertised-vs-measured are the same defect. This is the clause-(6) FALSE-MATCH at the metric layer — the instrument matched A number, not THE quantity — and a downstream consumer of an unverified metric INHERITS the bluff (§11.4.6: re-derive from the raw source before re-publishing; a figure carried between agents/briefs without its field-identity verification is a rumour with units).

(10) OBSERVER DECONTAMINATION — the instrument's own footprint is identified, excluded, and the decontaminated instrument re-proven seeing (extension, research-derived, 2026-07-23). Forensic FACT (genericised, 2026-07-22): a socket census read TIME-WAIT counts of 68–80 and nearly minted a leak finding — after excluding the SAMPLER'S OWN probe ports the real count was 4–5 (the census was ~16× observer-induced); the control needle proving the decontaminated instrument still SAW real state: the same scan did detect a genuine step-change (+264 MB RSS, 1 → 93 sockets) in a monitored process. Every measurement whose instrument CONTRIBUTES to the measured population (probe connections in a socket census, sampler CPU/IO in a load profile, probe processes in a process census, trace overhead in a latency measurement — the §12.12 self-match generalised from kill-lists to counts) MUST identify and EXCLUDE its own footprint before reporting, AND MUST prove via a (7)(b)-class control needle that the decontaminated instrument still sees genuine signal (an over-broad exclusion is a self-inflicted clause-(6) false-null). Composes §11.4.128(1) (observer-effect budget — the recorder must not perturb the recorded;

clause (10) adds the measurement-side half: where the footprint is unavoidable it is measured and SUBTRACTED, never reported as the system's own behaviour) and §11.4.119 (single-resource-owner — a measurement arm racing a concurrent driver of the same resource produces cross-contaminated evidence; partition or serialize FIRST, as measured work this cycle did by refusing to touch a sibling stream's in-flight file region).

(11) ARTIFACT-USABILITY, NOT PREREQUISITE-PRESENCE (extension, research-derived, 2026-07-23). Forensic FACT (genericised, 2026-07-22): an install/verify pipeline keyed its "can this host run the tool?" verdict on the presence of a BUILD prerequisite (a `command -v <compiler> probe`) — and turned a host holding a WORKING prebuilt binary into a false "cannot launch" FAIL, exit 1, in one transcript that showed BOTH "[ok] ... verified" AND "[FAILED] ... cannot launch" for the same artifact; the independent §11.4.142 review graded it BLOCKING (it breaks every upgrade on hosts that install from prebuilts). A readiness / install / health / build gate whose CLAIM is about an ARTIFACT ("the tool works here") MUST probe the ARTIFACT ITSELF — execute it through its real invocation path and assert its observable behaviour (§11.4.224(A)'s real-invocation-path discipline; §11.4.200's verify-the-target sibling at the readiness seam) — NEVER the presence of a prerequisite for ONE of several ways of producing it (a compiler, an SDK, a package manager). The proxy fails in BOTH directions at once: prerequisite-present ≠ artifact-usable (a toolchain with no artifact) AND prerequisite-absent ≠ artifact-unusable (a prebuilt with no toolchain) — clause (1) and clause (2) in a single check.

(12) SHELL-INSTRUMENT FOOTGUN CHECKLIST — a standing (7)(c) reference that shell-touching reviews MUST consult (extension, research-derived, 2026-07-23). The measured shell-instrument footgun classes are catalogued in the standing reference `docs/guides/shell_instrument_footguns.md` (inherited by reference per §11.4.28/§11.4.177, never copied per project). Among the founding, DEMONSTRATED classes: `timeout / nice / env -class EXEC WRAPPERS` cannot run shell functions/aliases (`rc=127 "command not found"` while the function is healthy — a clause-(1) false refusal; demonstrated 2026-07-23 under control conditions); `pgrep -f / pkill -f` matching the PROBING script's own command line (the §11.4.196(D)/§12.12 carrier — re-demonstrated 2026-07-23: a unique needle matched exactly ONE process, the probe's own harness); structural-delimiter text extractors (a `column-0- }` -terminated awk function extractor) that TRUNCATE the extracted unit

when its body legitimately contains the delimiter (heredoc/string content) and hand the test a PARTIAL artifact — a FALSE PASS when the truncated tail carried the code under test (demonstrated 2026-07-23: the tail marker vanished from the extraction); and xtrace (`bash -x`) BLINDED when the code under test redirects the shell's stderr (`exec ... 2>/dev/null` silenced the trace mid-stream in the real 2026-07-22 incident — countermeasure: route the trace via `BASH_XTRACEFD` to a dedicated fd the code under test does not own). Every §11.4.142/§11.4.209 code review of a shell-based test, gate, guard, or measurement MUST consult the checklist, and every zero/absence/PASS a shell instrument produces remains bound by (7)(b)'s class-matched control needle. The checklist is a LIVING document: a newly-measured footgun class is APPENDED with its captured demonstration (§11.4.5) in the same change-window, never carried as tribal memory (§11.4.215 — the binding reference lives tracked, in-repo).

Gate contract for clauses (9)-(12) (extension, 2026-07-23). `CM-GUARD-ASSERTS-REAL-CONDITION` is EXTENDED: **(iv)** every load-bearing metric cited in a gating / ranking / report context carries its verified field-identity + units + semantic class (capacity / occupancy / rate / cumulative / high-water) traced to the authoritative source definition — an unverified-semantics metric gating a decision → FAIL; **(v)** every census/profile whose instrument contributes to the measured population records its observer-exclusion AND a post-exclusion control-needle result — an undecontaminated census minting a finding → FAIL; **(vi)** every artifact-readiness verdict cites a probe of the artifact's own invocation path — a prerequisite-presence-only readiness verdict → FAIL; **(vii)** every review record touching a shell-based instrument cites the §11.4.201(12) checklist consultation. Paired §1.1 mutations: relabel a capacity field as occupancy in a gating metric → (iv) FAILs; strip the observer-exclusion from a census fixture → (v) FAILs; swap the artifact probe for a bare `command -v <compiler>` while the fixture host carries a working prebuilt → (vi) FAILs. Gate-code = a separate work item; this is the contract, not a claim that the code has shipped (§11.4.6).

Canonical authority: constitution submodule `Constitution.md` §11.4.201.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--proxy-signal-OK` , `--substring-match-guard` , `--false-positive-refusal-acceptable` , `--unvalidated-guard` , `--refuse-without-evidence` , `--zero-proves-absence` , `--skip-control-needle` , `--substring-is-good-enough` , `--carrier-match-OK` , `--trust-the-null` , `--metric-without-done-`

check , --gate-on-unvalidated-metric , --capacity-is-occupancy-OK , --skip-field-identity-check , --observer-footprint-included-OK , --prerequisite-presence-proves-artifact ,
--skip-shell-footgun-checklist flag.

§11.4.202 — Reporting directives: a report MUST auto-create a fully-populated, fully-synced workable item — never a prose acknowledgement (User mandate, 2026-07-15). Verbatim operator mandate: introduce reporting directives `ISSUE` / `BUG` / `TASK` that create a proper workable item filled with all details, with the WHOLE workable-items system automatically up to date and fully in sync — the main DB, all documents, and all external tracking systems; universal, reusable on every project, fully decoupled, recognized immediately by every project that pulls the constitution. Every governed project MUST expose REPORTING DIRECTIVES — the §11.4.140 registered actions `BUG` (Type=Bug), `TASK` (Type=Task), and `ISSUE` (generic report, CLASSIFIED into the §11.4.16 closed set) — such that a plain-language report from the operator is TURNED INTO a real, tracked, fully-synced workable item, automatically, every time. A report that is discussed, acknowledged, answered in prose, or "noted" WITHOUT landing a tracked item is a §11.4.202 violation of §11.4 PASS-bluff severity at the requirements-intake layer: the requirement was accepted and then silently evaporated (the §11.4.197 loss-of-requirements failure, at its entry point). **(1) GRAMMAR (§11.4.140).** The three directives are ordinary registry rows, so EVERY §11.4.140 form works out of the box (`BUG :: <text>` ,
`DEFAULT::BUG :: <text>` , `/DEFAULT::BUG <text>` , `BUG ---> <text>`), PLUS a SIXTH form this anchor adds — the **single-colon form** `NAME: <text>` (`ISSUE: subtitles are late`), the shape operators actually type. The single-colon form is **REGISTERED-ACTION-ONLY by construction**: a single-colon token that is NOT a registered action is a NO-OP (an ordinary prompt), NEVER an ASK and NEVER a sub-system shortcut — because ordinary English prose (`TODD:` , `WARNING:` , `FIXME:`) shares that shape (`NOTE:` is now a registered §11.4.140 severity marker), and routing it to the §11.4.66 clarify path would question every such sentence (§11.4.6 honest boundary: the other five forms keep their ASK-on-unknown behaviour — their shapes do not occur in prose). Host-command collisions are handled by the EXISTING §11.4.140 conflict mechanism (`slash_bare: auto + slash_conflicts: [..]`), which already satisfies the operator's explicit-prefix requirement: a bare `/NAME` that collides with a built-in (e.g. the documented Claude Code built-in `/bug`) is NOT honored, while the

namespaced `/DEFAULT::NAME` form is ALWAYS unambiguous — and the `NAME: / NAME :: / NAME --->` forms are never affected by any host command.

(2) THE ITEM. The created item MUST carry, on ALL surfaces: a valid Status (§11.4.15, `Queued` at intake) + a Type from the CLOSED set {Bug | Feature | Task} (§11.4.16 — `ISSUE` is a REPORTING CHANNEL, never a fourth type; inventing one is a violation) + a stable auto-incremented id (§11.4.54), and a COMPREHENSIVE structured description (§11.4.148 D2 / §11.4.171): what / affected-scope / reproduction / acceptance. An undetermined section is recorded as an explicit `UNKNOWN: gap` — NEVER invented (§11.4.6). For `ISSUE`, the type is classified from the report's content and stated as FACT; if the content does not determine it, the operator is ASKED (§11.4.66 / §11.4.105) BEFORE creation; ONLY when running autonomously where asking is impossible (§11.4.101) may the §11.4.16 lowest-stakes ambiguity default `Task` be used — and then the defaulted classification MUST be recorded verbatim in the item and surfaced for reclassification, never silently asserted.

(3) THE FULL SYNC (the load-bearing half). Creating the item is NOT sufficient: the directive MUST drive the WHOLE workable-items system into sync, in one shot — (a) the item lands in the SQLite single-source-of-truth (§11.4.93 / §11.4.95); (b) EVERY derived document is regenerated FROM the DB — the open + closed trackers, their summaries, and their `.md` / `.html` / `.pdf` / `.docx` siblings (§11.4.12 / §11.4.53 / §11.4.65 / §11.4.106); (c) the item is pushed to EVERY configured external tracker (§11.4.148 D5).

(4) NEVER A FAKED PUSH (§11.4.10 / §11.4.6 / §11.4.3). A tracker whose credentials or whose client are absent is SKIPPED with an honest machine-readable reason (`credentials_absent` — names of the unset variables ONLY, never a value; `tracker_client_absent` — PENDING-OPERATOR-INPUT). A tracker PASS may be recorded ONLY after a real push command really exited 0. Fabricating, assuming, or silently omitting a push is a §11.4 bluff at the integration layer; a missing tracker NEVER blocks the item from landing in the DB + docs (the report is never lost because a mirror is unreachable).

(5) DECOUPLED ENGINE, CONSUMER-OWNED DATA (§11.4.28 / §11.4.177). The item-creation + full-sync engine lives in the constitution submodule (`constitution/scripts/reporting/report_item.sh`), is inherited BY REFERENCE (never copied), and carries ZERO project literals; every project-specific value — DB path, id prefix, sync command, tracker commands + their required env vars — comes from a CONSUMER-OWNED config file (DATA, never an engine edit). The engine FAILS CLOSED with an actionable message when that config is absent — it never guesses a project's

paths (§11.4.6). **(6) ANTI-BLUFF (§11.4 / §11.4.107(10))**. The mechanism ships four-layer coverage: a pre-build gate (`CM-REPORTING-DIRECTIVES`) whose FUNCTIONAL invariant SOURCES and RUNS the parser (never a grep-only assertion, §11.4.108), a paired §1.1 mutation test that proves each invariant is genuinely load-bearing, and an end-to-end suite run against a REAL SQLite DB with REAL binaries (no mocks, §11.4.27) that asserts a real row lands with the right Type/Status/id, that the docs are regenerated FROM the DB, and that an absent tracker SKIPS honestly — self-validated with a golden-good + golden-bad + negative-control fixture set (an oracle that passes its golden-bad fixture is itself the bluff). Classification: universal (§11.4.17) — the consuming project supplies its DB path, id prefix, sync command, and tracker bindings per §11.4.35. Composes §11.4.140 (grammar) / §11.4.93 / §11.4.95 (DB SSoT) / §11.4.15 / §11.4.16 / §11.4.54 (status + type + id) / §11.4.148 / §11.4.171 (comprehensive description) / §11.4.106 / §11.4.12 / §11.4.53 / §11.4.65 (doc sync) / §11.4.10 (credentials) / §11.4.66 / §11.4.105 (clarify, never guess) / §11.4.101 (autonomous decision) / §11.4.197 (a started requirement never evaporates) / §11.4.28 / §11.4.177 (decoupling) / §11.4.164 (auto-registration on constitution pull) / §11.4.107(10) / §1.1. Propagation gate `CM-COVENANT-114-202-PROPAGATION` (literal `11.4.202`) + recommended gate `CM-REPORTING-DIRECTIVES` (registry declares ISSUE/BUG/TASK + the single-colon form is registered-only; the engine exists, is parse-clean, and is project-literal-free; the three directives really expand at runtime while prose does not; the honest-skip reasons exist and a tracker PASS is gated on a real exit 0) + paired §1.1 mutation (strip an action row, flip the single-colon form to non-registered-only, strip the parser branch, remove the prose NO-OP branch, make the engine fake a tracker push, or couple the engine to one project → the gate FAILs; strip the literal → the propagation gate FAILs).

Canonical authority: constitution submodule `Constitution.md` §11.4.202.

Non-compliance is a release blocker regardless of context. No escape hatch — no

`--report-without-item` , `--prose-acknowledgement-OK` , `--skip-item-creation` , `--skip-tracker-sync` , `--fake-tracker-push` , `--invent-a-fourth-type` , `--hardcode-project-in-engine` flag.

§ . . — Instant multi-stream resume engine: whole-fleet continuation state is a durable, content-addressed, atomically-committed snapshot resumed in $O(\text{changed})$ reads, not an $O(\text{history})$ re-read (User mandate, — —)

Compact summary: the continuation mechanism (§12.10 / §11.4.127 / §11.4.131 / §11.4.205) is EXTENDED — never replaced — by a mechanical engine so a fresh session resumes ALL work (every Track, main stream, and agent) at once, INSTANTLY and at token cost $O(\text{changed streams})$, by reading ONE snapshot manifest + one compact blob per stream, NEVER re-reading the full handoff history.

(1) **Content-addressed Merkle store** — each stream's canonical state bytes are named by their sha256 (dedup: unchanged streams share a hash → a global fleet snapshot costs $O(\text{changed})$); a snapshot is ONE manifest referencing every stream's content-hash (§11.4.205(4)). (2) **$O(\text{streams})$ resume, $O(\text{history})$ NEVER on the resume path** — `RestoreAll` reads the manifest + one blob per stream; the append-only event ledger is audit + lost-update recovery only, never read to resume (§11.4.127). (3) **Tamper-evident + deterministic** — every blob is integrity-checked on read (bytes MUST hash to their id) and a deterministic round-trip (re-serialize must byte-match), no wall-clock in the hashed content, verified by a SELF-VALIDATING oracle: golden-good PASS + golden-bad detected FAIL + negative-control (a legitimately-older-but-valid snapshot) PASS — the false-positive guard proving the verifier distinguishes a LAGGING copy from a TAMPERED one (§11.4.201/§11.4.107(10)/§11.4.206(3)). (4) **Durable + atomic + single-writer + crash-safe** — atomic ref write (temp → fsync → rename → dir-fsync, §11.4.205(6)), append-only ledger (§11.4.205(6)), single-writer-per-stream refusal (§11.4.206), advisory lock with provably-stale reap (`kill -0` liveness, never steals a live lock, §11.4.180/§9.2). (5) **Project-agnostic engine, consumer-owned data** — the engine (`github.com/vasic-digital/continuum` , GitLab mirror) carries ZERO project literals, fails closed rather than guess a store path (§11.4.6/§11.4.177), and is inherited BY REFERENCE as a depth-1 submodule under the §11.4.28(C) carve-out (`constitution/submodules/continuum/` , `helix-deps.yaml` , zero own-org deps); every consumer supplies its store path / actor / stream naming as DATA. Anti-bluff (§11.4/§11.4.108): the whole engine ships four-layer coverage (unit + integration + e2e + stress + chaos §11.4.85, race-clean) with a measured resume Metric (`blobs_read == 1 + streams` is the $O(\text{streams})$ proof) and the self-

validating oracle as captured evidence (§11.4.5/§11.4.69/§11.4.107); honest boundary (§11.4.6): it snapshots the state an agent REPORTS, not its execution image (not a CRIU/durable-execution runtime, §11.4.112), and `est_tokens` is a documented `bytes/4` heuristic, not a tokenizer count. Classification: universal (§11.4.17). Composes §12.10 / §11.4.127 / §11.4.131 / §11.4.205 / §11.4.116 / §11.4.147 / §11.4.180 / §11.4.201 / §11.4.206 / §11.4.107(10) / §11.4.28 / §11.4.177 / §11.4.85 / §9.2. Propagation gate `CM-COVENANT-114-207-PROPAGATION` (literal `11.4.207`) + recommended gate `CM-CONTINUUM-RESUME-ENGINE-PRESENT` (the engine present under `constitution/submodules/continuum/`, `go test -race ./... green, continuum selfcheck = good=PASS/bad=FAIL/negctrl=PASS`, zero project literals) + paired §1.1 mutation (strip the literal → propagation gate FAILs; downgrade the resume path to read the ledger, or the oracle to pass its golden-bad → the mechanism gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.207.

Non-compliance is a release blocker. No escape hatch — no `--resume-reads-history`^{11 4 208}, `--skip-integrity-check`, `--unvalidated-oracle-OK`, `--hand-typed-resume-state`, `--per-project-resume-reimpl` flag.

§ 11.4.208 – Operator-request-history document: every project maintains a project-local, always-in-sync ledger of every operator request/prompt (content + timestamp+timezone + track + alias + model + effort) (User mandate, – –)

Verbatim operator mandate (2026-07-15): introduce a request-history document that captures, per operator request/prompt — the full request content, the accepted timestamp with an explicit timezone, the track that processed it, the alias that took the workable item, and the model + effort used — always kept in sync, with no false data and no bluff of any kind. **The operator's load-bearing correction (§11.4.35 / §11.4.17):** the RULE (this universal discipline) belongs in the constitution submodule; the request-history DOCUMENT itself (project data) is PROJECT-LOCAL — the document is NEVER placed into the constitution.

Every governed project MUST maintain a single, project-local **operator-request-history document** — an append-only, newest-first ledger of EVERY operator request/prompt the project has received — at a fixed, project-declared standard path (declared once per §11.4.35, e.g. `docs/requests/history.md`, never moved without a §11.4.66 operator decision). Each entry MUST carry EXACTLY these operator-mandated fields: **(1) Request content** — the full/verbatim prompt text where archived, otherwise a faithful complete summary (each entry marks which — verbatim vs summary); **(2) Accepted (when)** — the exact date + time + TIMEZONE the request was accepted, with the timezone stated EXPLICITLY on every entry (the consuming project declares its default zone per §11.4.35; where only a date is durably recoverable the time is marked `time UNKNOWN`, never fabricated); **(3) Track** — the parallel-development track that processed it (the conductor/main track by default, per-domain tracks where applicable, §11.4.176/§11.4.178/§11.4.182); **(4) Alias** — the live agent/worker alias that took the workable item at acceptance (§11.4.182/§11.4.196); **(5) Model + effort** — the model + effort setting used. **(A) NO INVENTION (§11.4.6)**. A field not durably recoverable is recorded LITERALLY as `UNKNOWN` (with a one-line reason where useful) — NEVER guessed; a request-history entry that fabricates a timestamp/alias/model it never captured is a §11.4/§11.4.1 bluff at the requirements-record layer. **(B) HONEST RECONSTRUCTION BOUNDARY**. For pre-mandate sessions where per-prompt records were not archived, the document is best-effort reconstructed from the durable dated record (governance-mandate provenance dates, session ledgers, handoff docs, VCS history), with the unrecoverable per-prompt fields marked `UNKNOWN` and the reconstruction boundary stated explicitly in the document (§0-style scope note) — never silently implied complete. **(C) ALWAYS IN SYNC + EXPORTED**. The document carries the §11.4.44 revision header, is kept in sync with live state (§12.10 trigger set — every new operator prompt appends a row), and is exported to all mandated formats (§11.4.65/§11.4.73 — `.md` + `.html` + `.pdf` + `.docx` where the four-format mandate applies, each non-`.md` sibling's mtime ≥ the `.md`). **(D) KEEP-APPLYING MECHANISM**. The project ships an append mechanism (a helper script and/or a `UserPromptSubmit` -class hook) that appends a new newest-first row for each NEW operator prompt at capture time with the exact timestamp + track + alias + model + effort (track/alias derived deterministically per §11.4.182, honest `?` / `UNKNOWN` when undeterminable); an append helper that landed WITHOUT the automatic-capture hook is honestly a partial mechanism (the hook wiring is a tracked §11.4.197/§12.10 follow-up, never claimed as automatic capture it does not perform — §11.4.6). **(E) COMPLEMENTS**,

NEVER REPLACES. The request-history document complements — never replaces — any per-session no-loss requirements ledger (§11.4.197/§11.4.202), the workable-items DB SSoT (§11.4.93/§11.4.95), or the §12.10/§11.4.131 session-resumption artefacts; loss of a requirement at its intake (an operator prompt not recorded) is FORBIDDEN (§11.4.197). **(F) RULE-UNIVERSAL / DOCUMENT-PROJECT-LOCAL (§11.4.35).** This universal RULE lives here in the constitution submodule; the request-history DOCUMENT, its declared path, and the project's default timezone are consumer-owned project DATA and live ONLY in the consumer layer — the document is NEVER placed into the constitution submodule (§11.4.28 decoupling).

Classification: universal (§11.4.17) — a platform-neutral requirements-record discipline reusable by ANY project; the consuming project supplies its concrete document path, default timezone, and append-helper wiring per §11.4.35. Composes §11.4.6 (no-guessing — UNKNOWN never a fabricated field) / §11.4.35 (rule-universal / document-project-local split) / §11.4.44 (revision header) / §11.4.65 / §11.4.73 (four-format export) / §11.4.118 (enumerated-coverage honesty for the backfill) / §11.4.131 / §12.10 (always-in-sync live-state artefact) / §11.4.176 / §11.4.178 / §11.4.182 (track + alias derivation) / §11.4.196 (alias tracking) / §11.4.197 (no requirement lost) / §11.4.202 (report → workable-item — the request-history is the human-readable request ledger; §11.4.202 is its DB/tracker materialisation) / §11.4.93 / §11.4.95 (workable-items SSoT). Propagation gate CM-COVENANT-114-208-PROPAGATION (literal 11.4.208 across the consumer fleet) + recommended gate CM-REQUEST-HISTORY-DOC-PRESENT (the project-local request-history document exists at its declared path, carries the §11.4.44 revision header + the five mandated per-entry fields + an explicit-timezone stamp on every entry + a reconstruction-boundary scope note; the append helper exists + is parse-clean per §11.4.67; no fabricated field — UNKNOWN used for the unrecoverable) + paired §1.1 meta-test mutation (strip a mandated field from a row / drop the explicit-timezone stamp / fabricate an alias-or-model on an entry that never captured one / delete the document → the gate FAILS; strip the literal → the propagation gate FAILS; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md §11.4.208.

Non-compliance is a release blocker. No escape hatch — no `--request-history-optional`, `--skip-request-log`, `--no-timezone-stamp`, `--invent-missing-field`, `--request-history-in-constitution`, `--prompt-without-history-entry` flag.

2026 07 15

§ . . — Code-review MUST run on the Fable model at xhigh effort (Opus xhigh fallback) (User mandate, — —)

Verbatim operator mandate (2026-07-15): the mandatory independent code-review MUST ALWAYS be performed with the **Fable model at xhigh effort**; if Fable is unavailable, use **Opus at xhigh effort**.

Every mandatory independent code-review the constitution requires — the §11.4.125 code-review-agent gate before pre-build + main build, the §11.4.142 universal every-change-review (no exception), and the §11.4.194 exhaustive all-scenario/all-angle review + verify-against-captured-runtime-evidence, iterated to a zero-finding GO per §11.4.134 — MUST be dispatched on the **Fable model at xhigh effort**. **(A) Fallback (closed, ordered):** ONLY when Fable is genuinely unavailable (not configured / rate-limited / unreachable, established as FACT per §11.4.6 — never a guessed unavailability) the review MUST run on **Opus at xhigh effort**; a lower-effort or a non-Fable/non-Opus model is NOT an acceptable review substrate. xhigh effort is REQUIRED in both cases — the exhaustive §11.4.194 all-angle depth is not reachable at a lower effort. **(B) Independence preserved (§11.4.70/§11.4.20).** The Fable/Opus-xhigh model runs as the structurally-separated reviewer (never the author self-reviewing); §11.4.209 sets the MODEL + EFFORT of that review, it does not weaken the independence, the iterate-until-GO loop (§11.4.134), or the rock-solid captured-evidence requirement (§11.4.5/§11.4.69/§11.4.107). **(C) Honest boundary (§11.4.6).** A review claimed "done" that actually ran on a weaker model / lower effort than Fable-xhigh (or Opus-xhigh on genuine Fable-unavailability) is a §11.4 bluff at the review-substrate layer — the review's captured evidence MUST record which model + effort actually performed it (and, on fallback, the FACT of Fable's unavailability). STRENGTHENS §11.4.125/§11.4.142/§11.4.194/§11.4.134 — it does NOT replace any of them; it pins the review's execution substrate.

Classification: universal (§11.4.17). Composes §11.4.125 / §11.4.134 / §11.4.142 / §11.4.194 / §11.4.6 / §11.4.20 / §11.4.70 / §11.4.196 (alias/model selection).

Propagation gate `CM-COVENANT-114-209-PROPAGATION` (literal `11.4.209` across the consumer fleet) + recommended gate `CM-CODE-REVIEW-FABLE-XHIGH` (the mandatory code-review's captured evidence records model=Fable + effort=xhigh, OR model=Opus + effort=xhigh with a recorded Fable-unavailability fact) + paired §1.1 meta-test mutation (record a code-review run on a lower-effort or non-Fable/non-Opus model without a Fable-unavailability fact → the gate FAILs; strip the literal → the propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.209.

Non-compliance is a release blocker. No escape hatch — no

```
11 4 210
--review-any-model , --review-lower-effort-OK , --skip-fable , --non-
xhigh-review , --fallback-without-fable-unavailability-proof flag.
```

§ . 2026 07 15 — Zero-loss request/prompt intake: every operator request MUST be mechanically captured, tracked, and processed — never skipped, ignored, avoided, or lost (User mandate, — —)

Forensic anchor — verbatim user mandate (2026-07-15): "We MUST CHECK which requests / prompts we have given havent passed through except the accumulative prompts / requests history document !!! There may be more of it! This MUST NEVER happen! We MUST make sure that every request / prompt we issues is AWLAYS respected and taken into the account, executed and processed! This is MANDATORY RULE !!! There MUST NOT BE any avoiding, ignoring, skipping or any form of loss for requests / prompts we make! Extend properly constitution Submodule and add mandatory linkings, hooks and extend with dcos chain any part of the System we MUST in order for this to be permanently resolved! Applying this starts immidiately! Commit and push all changes to constitution Submodule to all upstreams!"

Forensic FACT (2026-07-15): an entire session's worth of operator prompts (a release-drive "keep going", a `BACKGROUND` emulator directive, a live-testing model matrix, an alias switch, a multi-track directive) went UN-RECORDED in the §11.4.208 request-history ledger because §11.4.208(D)'s "keep-applying

mechanism" (the `UserPromptSubmit` auto-capture hook) was left as a "tracked follow-up / honestly partial" — so the ledger drifted the moment hand-editing stopped, exactly the §11.4.205 failure class ("a rule documented as mechanically enforced but implemented nowhere is worse than no rule").

Every operator request / prompt MUST be (a) **CAPTURED** — recorded in the §11.4.208 request-history ledger MECHANICALLY at intake, never depending on manual vigilance; (b) **TRACKED** — mapped to a tracked action / workable item (§11.4.202 report → item, §11.4.197 completion) OR an explicit evidence-backed no-op-with-reason; and (c) **PROCESSED** — respected, acted upon, executed, never skipped / ignored / avoided / dropped. Loss of a request at ANY of these three stages is a §11.4 PASS-bluff at the requirements-intake layer of the SAME severity class as §11.4.197 (loss of requirements FORBIDDEN) — §11.4.210 is the INTAKE-side MECHANICAL guarantee that §11.4.197's no-loss holds from the very first moment a request arrives. §11.4.210 STRENGTHENS §11.4.208(D): the auto-capture hook is PROMOTED from "tracked follow-up / honestly partial" to MANDATORY, wired, and verified.

Mandatory protections (ALL hold): (1) **MECHANICAL CAPTURE, not manual vigilance (§11.4.205).** The §11.4.208 ledger's keep-applying mechanism MUST be a WIRED, VERIFIED-RUNTIME-STATE auto-capture hook (a `UserPromptSubmit` -class hook that appends every operator prompt to the ledger at intake) — NOT a "tracked follow-up", NOT a manual conductor step. "Installed" = a verified runtime state observed in the live hook path (§11.4.205(3)), never a doc claim; the hook MUST be non-fatal (never blocks / fails a prompt) and honest (`UNKNOWN` fields never invented, §11.4.6). (2) **DOCS-CHAIN BOUND (§11.4.106).** The ledger + its four-format exports (§11.4.65/§11.4.73) are bound into the docs-chain so they stay in sync mechanically; a stale ledger export is a §11.4.106 sync violation. (3) **GATED (§11.4.75, enforced-not-advisory).** A pre-build gate FAILs if the auto-capture hook is not wired in the live hook path (verified-installed per §11.4.205(3), content-matched to its tracked source), OR the ledger is stale relative to the durable prompt record. Paired §1.1 mutation — un-wire the hook / orphan the ledger → gate FAILs. (4) **AUDIT + NO SILENT DROP.** Any operator prompt found un-captured (in `.remember/` logs, older CONTINUATION revisions, session transcripts) MUST be back-filled into the ledger, never silently dropped (§11.4.197); every captured request resolves to a tracked action / item (§11.4.202) OR an explicit evidence-backed no-op-with-reason — captured-then-ignored is forbidden. (5) **PROCESS EVERY CAPTURED REQUEST.** The autonomous loop (§11.4.126 /

§11.4.87 / §11.4.94 / §11.4.97 / §11.4.103) treats every open captured request as an actionable queue item; the loop's done-condition MUST NOT read satisfied while any captured request is un-processed (mirrors §11.4.147(e)). (6) **HONEST BOUNDARY (§11.4.6)**. Machine-derivable fields (timestamp / track / alias) are captured verbatim; model / effort where undeterminable are UNKNOWN , never invented; a back-filled catch-up row honestly marks its reconstruction boundary (§11.4.208(B)) — never a fabricated time.

Classification: universal (§11.4.17) — the consuming project supplies its ledger path, hook path, and docs-chain context per §11.4.35. Composes §11.4.197 (§11.4.210 is its intake-side mechanical enforcement) / §11.4.202 (report → tracked item) / §11.4.208 (the ledger document — §11.4.210 promotes its (D) hook from partial-follow-up to mandatory) / §11.4.205 (enforced-not-advisory, hook-installed-is-verified-runtime-state) / §11.4.106 (docs-chain) / §11.4.126 / §11.4.87 / §11.4.94 / §11.4.97 / §11.4.103 (the loop processes every captured request) / §11.4.6 / §11.4.75 / §11.4.147(e). Propagation gate CM-COVENANT-114-210-PROPAGATION (literal 11.4.210) + recommended gate CM-REQUEST-CAPTURE-HOOK-WIRED (the auto-capture hook is verified-installed in the live hook path AND content-matches its tracked source per §11.4.205(3); the ledger is fresh; docs-chain-bound) + paired §1.1 meta-test mutation (un-wire the hook → recommended gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md §11.4.210.

Non-compliance is a release blocker. No escape hatch — no --skip-request-capture^{11 4 211} , --manual-ledger-OK , --hook-optional , --drop-unprocessed-request , --ledger-may-be-stale , --request-may-be-skipped flag.

§^{26 07 16} . . — Merge-conflict resolution during main→feature/product/flavor merges MUST run on the Fable model at xhigh effort (Opus xhigh fallback) (User mandate, — —)

Verbatim operator mandate (2026-07-16): "During the merge of main branch to all feature/product/flavor branches of tracks 2 - 4, if conflicts happen they MUST BE resolved using Fable model with xhigh effort! If Fable is not available, then with

Opus with xhigh! Add this MANDATORY RULE into constitution Submodule and its files CONSTITUTION.md, CLAUDE.md, AGENTS.md, QWEN.md, GEMINI.md, and other related relevant files!"

Any merge-CONFLICT resolution performed while integrating the canonical trunk (`main / master`) into a feature / product / flavor branch — the §11.4.188 regular main → feature merge cadence, the §11.4.195 branch-taxonomy integration, the §11.4.41 merge-first pipeline, the §11.4.167(D) trunk-into-stream sync AND the §11.4.167(I) approval-gated back-merge — MUST be carried out on the **Fable model at xhigh effort**. This is the merge-resolution analogue of §11.4.209 (which pins the same model + effort for code-REVIEW): §11.4.211 pins it for the DECISION-MAKING of conflict resolution (which side to keep, how to UNION both sides without loss, marker removal, semantic reconciliation of divergent edits). **(A) Fallback (closed, ordered):** ONLY when Fable is genuinely unavailable (not configured / rate-limited / unreachable, established as FACT per §11.4.6 — never a guessed unavailability) the conflict resolution MUST run on **Opus at xhigh effort**; a lower-effort or a non-Fable/non-Opus model is NOT an acceptable substrate for resolving merge conflicts. `xhigh effort` is REQUIRED in both cases — a careless conflict resolution silently drops a commit or commits a marker (§9 / §11.4.41 step 3), so the highest-effort scrutiny is mandatory. **(B) Scope.** The operator named tracks 2–4 (`feature / product / flavor` branches); the rule GENERALISES per §11.4.17 to EVERY controlled main → (feature|product|flavor) merge on EVERY track that hits a conflict — the substrate requirement does not depend on which track. **(C) Safety preserved.** §11.4.211 sets the MODEL + EFFORT of the conflict-resolution work; it does NOT weaken and is BOUNDED BY the surrounding merge-safety invariants — §11.4.113 (NEVER force-push), §9 / §9.2 (no commit lost, hardlinked pre-op backup before any large/risky merge), §11.4.41 step 3 (ZERO conflict markers committed, ZERO file silently dropped, union of both sides preserved), §11.4.188(4) anti-bluff (post-merge smoke GREEN + conflict-marker scan empty + no-lost-commit, captured evidence §11.4.5/§11.4.69), §11.4.84 (quiescent-only), §11.4.37/§11.4.71 (fetch-first). **(D) Honest boundary (§11.4.6).** A merge whose conflicts were resolved on a weaker model / lower effort than Fable-xhigh (or Opus-xhigh on genuine Fable-unavailability) is a §11.4 bluff at the merge-resolution-substrate layer — the merge's captured evidence MUST record which model + effort resolved the conflicts (and, on fallback, the FACT of Fable's unavailability). STRICT EXTENSION of §11.4.209 (which pins Fable-xhigh / Opus-xhigh for code-REVIEW) applied to merge-CONFLICT resolution; STRENGTHENS

§11.4.188 / §11.4.195 / §11.4.41 by pinning the substrate of their conflict-resolution step — the two anchors are the review-side (§11.4.209) and merge-side (§11.4.211) halves of the same Fable-xhigh substrate discipline.

Classification: universal (§11.4.17). Composes §11.4.209 / §11.4.188 / §11.4.195 / §11.4.41 / §11.4.113 / §9 / §9.2 / §11.4.6 / §11.4.167(D)(I) / §11.4.84 / §11.4.37 / §11.4.71 / §11.4.196 (model/alias selection). Propagation gate `CM-COVENANT-114-211-PROPAGATION` (literal `11.4.211` across the consumer fleet) + recommended gate `CM-MERGE-CONFLICT-FABLE-XHIGH` (a main → feature/product/flavor merge that resolved conflicts records model=Fable + effort=xhigh, OR model=Opus + effort=xhigh with a recorded Fable-unavailability fact) + paired §1.1 meta-test mutation (record a conflict-resolving merge on a lower-effort or non-Fable/non-Opus model without a Fable-unavailability fact → the gate FAILs; strip the literal → the propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.211.

Non-compliance is a release blocker. No escape hatch — no `--merge-conflict-any-model`, `--skip-fable-merge-resolution`, `--non-xhigh-merge-conflict`, `--resolve-conflict-lower-effort-OK`, `--fallback-without-fable-unavailability-proof` flag.

§_{26 07 16} . . . — Main README is the canonical starting-point / entry point for ALL project documentation — no doc may be an orphan unreachable from README (User mandate, — —)

Verbatim operator mandate (2026-07-16): "Manual tests catalog must have linking all through the main README as the starting point of this and any other Project documentation! This fact and mandatory rule MUST be added into the constitution."

The main README (`README.md` at the repository root) is the canonical STARTING POINT / entry point for ALL project documentation, for THIS project and EVERY project incorporating this constitution. Every project document in the §11.4.65 export scope — INCLUDING the manual-tests catalog (the §11.4.185 manual-QA testing catalog), the §11.4.153 per-feature Status / Status_Summary docs, the Issues / Fixed trackers + their summaries (§11.4.12 / §11.4.53 / §11.4.57),

and every guide / research note / changelog / script companion doc — MUST be reachable / linked (DIRECTLY or TRANSITIVELY) from the main README. No project document may be an ORPHAN — a doc that no link-path from README reaches is a §11.4.212 violation. This is the STRICT GENERALISATION of §11.4.57: §11.4.57 mandated a `Tracked-Items + Status Documents` doc-link section in the README covering the tracker / status doc set; §11.4.212 extends the same reachability guarantee to EVERY §11.4.65-scope document, so the README is the single navigable root of the project's documentation tree. **(A) Reachability = a link path.** README → guide → sub-doc is fine (transitive); the requirement is that SOME path from README reaches every in-scope doc, not that every doc is linked directly at top level. A doc no path reaches is an orphan. **(B) Always-sync (§11.4.59 / §11.4.106 docs-chain).** The README doc-graph is kept current mechanically — a newly-added §11.4.65-scope doc that no README-reachable path links is a §11.4.212 violation the moment it lands; the manual-tests catalog specifically MUST be linked from the README (the operator's load-bearing example). **(C) Honest boundary (§11.4.6).** Reachability is a link-graph property (every doc is FINDABLE from README), NOT a claim about each doc's CORRECTNESS or freshness (those stay §11.4.44 / §11.4.106 / the doc's own gates). **(D) Composition with §11.4.57.** §11.4.57's `Tracked-Items` section remains the mandated README landing block for the tracker / status set; §11.4.212 adds that EVERY OTHER in-scope doc is likewise reachable — §11.4.57 is the specialisation, §11.4.212 the generalisation.

(E) README INTRODUCTION + ILLUSTRATION (extension, 2026-08-15). The main README MUST open with a SHORT introduction (what the project is, who it is for, its current state at a glance — the reader's 30-second answer) AND MUST ILLUSTRATE all work done: for every major capability the project ships, the README carries either a direct summary paragraph OR a link to a detailed document explaining WHAT was done AND WHY it was done (never orphaned links, never "see the changelog" as the only explanation). Every linked document MUST justify its own existence in one sentence at the link site — a link with no "why this matters" is a §11.4.212(E) violation of §11.4.91 summary-clarity severity at the entry-point layer.

(F) LINKAGE-COMPLETENESS RATCHET (extension, 2026-08-15). The reachability graph from README (transitively per (A)) MUST cover the full §11.4.65-scope doc-set enumeration with a monotone-decreasing UNREACHED count (mirrors the §11.4.135 ratchet pattern); a §11.4.65-scope doc that is not

reachable from README at pre-build/release seams FAILs the extended `CM-README-DOC-ENTRYPOINT-COMPLETE` gate; brownfield adoption per operator §11.4.66 (§11.4.224(E) fence). The existing `CM-COVENANT-114-212-PROPAGATION` gate stays unchanged (literal `11.4.212` unchanged, propagation stays GREEN); `CM-README-DOC-ENTRYPOINT-COMPLETE` is EXTENDED with clauses (E)/(F); a new paired mutation strips the introduction, OR removes a "why this matters" annotation, OR breaks reachability of an in-scope doc → the extended gate FAILs.

Classification: universal (§11.4.17) — the consuming project supplies its README path, its manual-tests-catalog path, and its §11.4.65 doc-set enumeration per §11.4.35. Composes §11.4.57 / §11.4.59 / §11.4.65 / §11.4.106 / §11.4.153 / §11.4.185 / §11.4.44 / §11.4.12 / §11.4.53. Propagation gate `CM-COVENANT-114-212-PROPAGATION` (literal `11.4.212` across the consumer fleet) + recommended gate `CM-README-DOC-ENTRYPOINT-COMPLETE` (every §11.4.65-scope document is reachable — directly or transitively — from the main README; the manual-tests catalog is linked) + paired §1.1 meta-test mutation (orphan an in-scope doc so no README path reaches it → the gate FAILs; strip the literal → the propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.212.

Non-compliance is a release blocker. No escape hatch — no `--orphan-doc-OK`, `--readme-not-entrypoint`, `--skip-readme-link`, `--doc-unreachable-OK` flag.

§ 11.4.212 — **FEATURE research-scheduling directive: recognized via all § 11.4.212 forms, SCHEDULES (never synchronously executes) a deep, enterprise-grade research + implementation-planning effort as a tracked workable item (operator-directed mandate, — —)**

Mandate (2026-07-16): extend the §11.4.202 reporting-directive family with a fourth registered action, `FEATURE`, that turns a plain-language feature description into a SCHEDULED, deep, enterprise-grade research-and-planning effort — mirroring EXACTLY how `ISSUE` / `BUG` / `TASK` turn a report into a tracked workable item, but for RESEARCH SCHEDULING rather than defect/task/feature reporting: the directive itself creates the tracked item and enqueues it durably; the multi-day

research/planning/documentation work it describes is EXECUTED LATER by the project's standing autonomous loop when it claims that item, never synchronously by the directive engine.

FEATURE is a research-SCHEDULING directive, structurally parallel to the §11.4.202 report-creation directives: recognized via ALL §11.4.140 forms —

FEATURE :: <desc> (bare ::), DEFAULT::FEATURE :: <desc> (namespaced ::), /FEATURE <desc> (bare slash), /DEFAULT::FEATURE <desc> (namespaced slash), FEATURE ---> <desc> (arrow) — PLUS the §11.4.202 single-colon form FEATURE: <desc> , REGISTERED-ACTION-ONLY by construction exactly like BUG: / TASK: / ISSUE: : because the action token grammar is UPPERCASE-only ([A-Z][A-Z0-9_]*), ordinary lowercase prose such as "feature: add dark mode" never matches ANY of the six forms and is never expanded and never asked about (§11.4.6) — only the literal uppercase FEATURE token triggers the directive. \FEATURE (leading backslash) escapes to a literal per the existing §11.4.140 escape mechanism, unchanged.

(1) Scheduling, not synchronous execution — the load-bearing distinction

from every other §11.4.140 action. BACKGROUND / CRITICAL / IMPORTANT / NOTE tag HOW the remainder of the SAME turn is handled; BUG / TASK / ISSUE immediately create + fully sync a workable item describing a report and the reported matter is then handled inline. FEATURE differs: its OWN execution is fast and bounded — create the tracked item (Type=Task, Status=Queued, §11.4.54 stable id), append it to a durable, never-dropped feature-research queue, sync the DB-derived documents, and push to configured trackers — exactly as §11.4.202 does for TASK , and via the SAME engine, never a duplicate implementation. The deep multi-day research / systematic-debug / documentation / planning work DESCRIBED INSIDE the item is NOT performed by the directive at dispatch time; it is performed LATER by the project's standing autonomous loop (§11.4.87 endless-loop / §11.4.94 zero-idle / §11.4.97 maximum-idle-use / §11.4.103 continuous-parallel-stream / §11.4.126 default-autonomous-loop) the moment it claims the item, and is driven to a genuinely COMPLETED (implemented AND wired AND verified) or explicitly evidence-backed CLOSED terminal state under the §11.4.197 research/kicked-off-work completion mandate — a FEATURE item is NEVER permitted to sit un-wired in the backlog; "scheduled" is not "done" and is never reported as such (§11.4.6).

(2) The scheduled research-and-planning mandate — embedded verbatim into the created item's comprehensive description (§11.4.148 D2 / §11.4.171) as its acceptance-defining WORK PROGRAM: (a) a deep web-research series — articles, guides, scientific papers, open-source projects/codebases, worked examples — on the best way to incorporate the described feature into the project (§11.4.8 deep-research-before-implementation / §11.4.99 latest-source cross-reference / §11.4.150 deep multi-angle research per change); (b) systematic-debug (§11.4.102) of ALL data obtained to enumerate EVERY weak spot / gap / danger-zone / inconsistency / imperfection the research surfaces, and design a risk-free, rock-solid solution for EACH one found — no gap left unaddressed; (c) the design MUST ALWAYS be enterprise-grade, bleeding-edge, and innovative — never less; (d) MANDATORY exhaustive documentation — technical documents, an implementation plan down to lines-of-code + micro-proof-of-concepts, diagrams / schemes / graphs, illustrations, SQL definitions, templates, and every other relevant material, four-format-exported and revision-headed per §11.4.65 / §11.4.73; (e) investigate + plan REUSE of existing `vasic-digital` + `HelixDevelopment` submodules/components BEFORE proposing a rewrite, extending them freely where genuinely needed while keeping them fully decoupled / project-not-aware / reusable per §11.4.28 / §11.4.74 / §11.4.177; (f) plan from the TESTING point of view from day one — every supported test type, the Challenges submodule, and full HelixQA bank coverage per §11.4.27 / §11.4.169; (g) divide the work into phases / tasks / subtasks, fine-grained and nano-detailed enough to drive integration / implementation / wiring / testing / scaling directly, no vague placeholders; (h) plan explicitly for enterprise scalability and maximal performance; (i) when the feature has a UI/UX surface, produce full wireframes / diagrams / design files (Figma / PSD / PDF, or whatever formats the project mandates) authored via OpenDesign per §11.4.162 / §11.4.190; (j) plan full CodeGraph integration with regular / real-time index synchronisation per §11.4.78 / §11.4.79 / §11.4.80; (k) create fully-detailed workable items (every detail + reference + attachment) synced in REAL TIME to the SQLite single-source-of-truth (§11.4.93 / §11.4.95) + every derived workable-items document + every related project doc/ component + every connected external work-tracking system (ClickUp, HelixTrack, JIRA, or any other configured tracker per §11.4.148 D5) — honestly SKIPPING any absent tracker with a machine-readable reason (`credentials_absent` / `tracker_client_absent` , §11.4.10 / §11.4.6) rather than EVER faking a push, exactly as §11.4.202 mandates.

(3) The item. The created item carries Type=Task (§11.4.16 — a scheduled research/planning effort is itself an internal workstream, not a product defect nor yet a shipped end-user capability; the feature it researches may spawn its OWN Feature-typed implementation item(s) once planned) + Status=Queued (§11.4.15) + a stable auto-incremented id (§11.4.54) + the §11.4.148 / §11.4.171 comprehensive structured description embedding (i) the verbatim feature description, (ii) the full research/planning/testing/scalability/UI/CodeGraph/tracker mandate of clause (2) verbatim as the item's WORK PROGRAM, and (iii) acceptance criteria = the full research-doc tree is produced under the item's recorded destination path, every enumerated gap carries a designed solution, an implementation plan down to lines-of-code exists, the planned follow-on workable items are created, and the whole effort is validated per §11.4.197 (never left un-wired). An undetermined section is recorded as an explicit UNKNOWN: gap, never invented (§11.4.6), exactly as §11.4.202's engine already does for BUG/TASK/ISSUE.

(4) Decoupled engine, reused machinery, consumer-owned data. The item-creation + full-sync half of the FEATURE mandate is NOT reimplemented: the engine (`constitution/scripts/feature/schedule_feature_research.sh`) invokes the SAME §11.4.202 `report_item.sh` engine to create the item, sync the derived documents, and push to trackers — so the already-audited no-faked-push / decoupling / anti-bluff invariants of §11.4.202 are inherited rather than duplicated (a second, divergent implementation of tracker-push logic would itself be a bluff-surface regression). The FEATURE-specific additions — the research-doc-tree destination path and the durable `docs/requests/feature_queue.md` queue entry (mirroring the §11.4.140 BACKGROUND action's `docs/requests/background_queue.md` durability pattern, so a scheduled request is NEVER silently dropped or forgotten) — are the engine's OWN, narrow responsibility. Both engines are inherited BY REFERENCE, carry ZERO project literals, and fail closed with an actionable message when their consumer config is absent (§11.4.6 / §11.4.28 / §11.4.177); every project-specific value (DB path, id prefix, research-doc root, sync command, tracker bindings) is consumer-owned DATA in `feature.yaml` , never an edit to the engine.

(5) Anti-bluff (§11.4 / §11.4.107(10)). Four-layer coverage mirroring §11.4.202: a pre-build gate (`CM-FEATURE-DIRECTIVE`) whose FUNCTIONAL invariant SOURCES and RUNS the §11.4.140 parser (never a grep-only assertion, §11.4.108) to prove the FEATURE token genuinely expands via the single-colon, `::` , arrow, and slash forms while lowercase prose (`feature: ...`) does NOT

expand and is NEVER asked about; a paired §1.1 mutation test proving every invariant (registry presence, engine parseability, decoupling, the durable-queue append, delegation-not-reimplementation of the tracker machinery, and the scheduling-not-synchronous framing) is genuinely load-bearing; and captured evidence of a real feature-research item landing in a temp SQLite DB, its derived documents regenerated FROM the DB, and its trackers SKIPping honestly.

Classification: universal (§11.4.17) — the consuming project supplies its DB path, id prefix, research-doc root, sync command, and tracker bindings per §11.4.35.

Composes §11.4.140 (grammar) / §11.4.202 (the sibling reporting-directive family + its reused engine) / §11.4.197 (a started research effort never evaporates) / §11.4.8 / §11.4.99 / §11.4.150 (deep research) / §11.4.102 (systematic-debug of every gap) / §11.4.65 / §11.4.73 (exhaustive documentation) / §11.4.28 / §11.4.74 / §11.4.177 (reuse-first, decoupled) / §11.4.27 / §11.4.169 (test-type + Challenges + HelixQA planning) / §11.4.162 / §11.4.190 (OpenDesign UI planning) / §11.4.78 / §11.4.79 / §11.4.80 (CodeGraph integration) / §11.4.93 / §11.4.95 / §11.4.148 (workable-items + tracker sync) / §11.4.54 / §11.4.15 / §11.4.16 (id/status/type) / §11.4.10 (credentials) / §11.4.6 (no-guessing) / §11.4.87 / §11.4.94 / §11.4.97 / §11.4.103 / §11.4.126 (the autonomous loop that actually performs the scheduled work). Propagation gate `CM-COVENANT-114-213-PROPAGATION` (literal `11.4.213` across the consumer fleet) + recommended gate `CM-FEATURE-DIRECTIVE` (registry declares `FEATURE` + its single-colon form is registered-only; the engine exists, is parse-clean, and is project-literal-free; the `FEATURE` directive really expands at runtime via all six forms while lowercase prose does NOT; the engine SCHEDULES via the durable feature-research queue AND delegates item-creation/sync/tracker-push to the §11.4.202 `report_item.sh` engine rather than reimplementing it) + paired §1.1 meta-test mutation (strip the `FEATURE` registry row, OR flip its single-colon form to non-registered, OR strip the parser branch, OR remove the durable-queue append, OR couple the engine to one project, OR make the engine reimplement tracker logic instead of delegating, OR strip the scheduling-not-synchronous framing from the expansion text → the gate FAILs; strip the literal → the propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.213.

Non-compliance is a release blocker. No escape hatch — no `--feature-runs-synchronously` , `--skip-feature-queue` , `--feature-without-research-mandate` , `--reimplement-tracker-push` , `--hardcode-project-in-engine` , `--skip-item-creation` flag.

2026 07 17

§ . . — **Recurrence-links-not-mints:**
a defect that returns **MUST** reopen its existing
item, never enter as a new id (research-derived,
- -)

Forensic anchor (genericised, 2026-07-17). In a consuming project, operator-reported defects that were RECURRENCES of already-tracked, already-"fixed" defects re-entered the tracker as BRAND-NEW ids rather than reopening the originals (measured: at least five distinct recurrence chains, one three deep). A recurrence filed as a new id **never increments the original's reopen counter** — so the §11.4.55 `reopens_count` , which is the exact signal §11.4.132(d) and §11.4.189 use to rank the most-fragile work for the deepest scrutiny, **stays quiet precisely on the items that keep breaking**. The project's recorded reopen-per-fix rate was 52% (published elite <10%) and was simultaneously a KNOWN UNDERCOUNT for this reason: the metric that was supposed to find the fragile items was silenced by the intake path. The same defect inflates the queue and forks the work — in one measured window ~15 tracker items mapped to ~9 real root causes, so "fix all 15" would have produced duplicate and mutually-conflicting fixes on one cause.

The mandate. Before minting a NEW id for any incoming defect — from ANY intake path — the system **MUST** check whether an existing item already describes the SAME defect, and if so LINK the report to that item rather than mint a new id, and REOPEN that item **if and only if it is currently in a terminal state** (§11.4.33 `Fixed` / `Implemented` / `Completed` , or a §11.4.90 closure OTHER than `duplicate-of` — see the chain rule).

(0) RESOLVE TO THE CANONICAL CHAIN HEAD FIRST — a `duplicate-of` closure is a POINTER, not a destination. §11.4.90's closure vocabulary includes `duplicate-of` , and clause (5) below MANUFACTURES exactly such items; the (2) key will match them, because a duplicate shares its original's subject BY CONSTRUCTION. Therefore, before applying link-and-iff-terminal-reopen, every match **MUST** be resolved THROUGH its `duplicate-of` links to the CANONICAL

CHAIN HEAD (the item the chain ultimately points to), and the link + any reopen land on THAT head — never on an intermediate duplicate. Reopening a `duplicate-of` closure would resurrect a non-canonical copy alongside its still-live canonical item: two live items for one defect — the exact queue-inflation + forked-fixes disease this anchor's own forensic names — and it would credit the §11.4.55 reopens signal to the WRONG item, defeating clause (6), which IS the point, on precisely the chain-class items the anchor was minted for. **Multi-match resolution:** when the key matches several items (an original plus its duplicates, possibly plus an open item), resolve every match to its chain head, then act on the SINGLE surviving head; if that resolution yields more than one distinct head the verdict is UNDECIDED per clause (3) (ask, or mint-with-link — never guess which head is meant, §11.4.6); if the resolved head is OPEN, prefer it and LINK ONLY. A chain whose head cannot be resolved (a broken/circular `duplicate-of` link) is UNDECIDED + a tracked integrity defect (§11.4.186 orphan-ref class), never a silent pick.

Open-original case (the ordinary two-reports-of-one-live-bug): when the resolved chain head is still OPEN — any non-terminal §11.4.15 status, i.e. `Queued` / `In progress` / `Ready for testing` / `In testing` / `already Reopened` / `Operator-blocked` (§11.4.21's 7th value, non-terminal) — the correct action is LINK ONLY — "reopen an open item" is undefined in the §11.4.15 status machine, and §11.4.34's `Reopened` status + its reason vocabulary presuppose a prior closure to demote FROM (§11.4.7); forcing a reopen there would corrupt both the status and the counter. The reopen-and-increment path exists to correct the specific lie "this was closed as fixed and it is not" — an already-open item is not telling that lie.

§11.4.202 precedence (stated explicitly — never left for a consumer to reconcile). §11.4.202 mandates that every report lands a tracked item, "automatically, every time", and forbids `--skip-item-creation` ; §11.4.214 forbids `--mint-every-report` . These are NOT in conflict and MUST NOT be read as such: §11.4.202's requirement is that **no report is ever LOST** — it is satisfied IN FULL by a SAME-DEFECT report landing as a link + (terminal-only) reopen on the existing item, because the report is tracked, the item is live, and the §11.4.55 counter now tells the truth. §11.4.202 forbids the report evaporating into prose; it does NOT require a fresh id per report. §11.4.214 refines WHICH item a report lands on; §11.4.202 owns THAT it lands. Where the check verdict is DISTINCT or UNDECIDED, §11.4.202's mint-a-new-id path runs unchanged. **(1)**

EVERY INTAKE PATH IS BOUND, not merely the reporting channel. Defects enter via the §11.4.202 reporting directives, via AI-detected test/gate failures, via regression-guard verdicts, via cycle re-discovery, via captured-evidence retrospect, and via operator manual testing — §11.4.34's own reason vocabulary (`test-failed` , `cycle-re-discovered` , `manual-testing-detected` , `end-user-report`) enumerates them. A dedup check wired ONLY into the reporting channel under-covers exactly the machine-driven paths that recur most. **(2) THE KEY IS STRUCTURED, NEVER A BARE SUBSTRING (§11.4.186).** Match on ticket reference where present, else on a normalised (`subject` , `scope`) — explicitly NOT a bare subject substring, which is the false-merge generator. **(3) A WRONG MERGE SILENTLY DELETES A REAL DEFECT — so `distinct-but-similar` is a FIRST-CLASS VERDICT, not a failure.** The check has three outcomes: SAME-DEFECT (link + reopen), DISTINCT (mint a new id — correct and expected), and UNDECIDED (§11.4.6: do NOT guess — ask per §11.4.66/§11.4.105, or mint the new id WITH a recorded candidate-duplicate link for review; when autonomous per §11.4.101, prefer minting-with-link over merging, because a spurious extra id is cheap and recoverable while a wrongly-merged defect is silently LOST — the asymmetry decides the default). **(4) THE DEDUP ORACLE IS SELF-VALIDATED (§11.4.107(10) / §11.4.186).** It ships a golden-good (a true recurrence — MUST link+reopen), a golden-bad (a missed recurrence — MUST be caught), and a **negative-control**: two genuinely DISTINCT items that merely share a subject, which MUST NOT be merged. An oracle without the negative control is a false-merge engine and is itself a §11.4 bluff. **(5) ID STABILITY IS PRESERVED (§11.4.54).** Linking a recurrence to its original NEVER rennumbers, reuses, or retires an id: the original is REOPENED and its counter increments; an already-minted duplicate is LINKED to the original (a recorded duplicate-of relation per §11.4.90) — never silently deleted, never renumbered. **(6) THE RESTORED SIGNAL IS THE POINT.** Reopening the original is what makes §11.4.55's counter true, which is what makes §11.4.132(d)/§11.4.189 rank the genuinely-fragile items first. A reopen counter that undercounts is not a cosmetic defect — it is a §11.4/§11.4.1 bluff at the prioritisation layer: it reports the fragile set as stable, and the deepest scrutiny is then aimed at the wrong items.

Honest boundary (§11.4.6). Dedup-at-intake makes the reopen signal TRUE; it does not make the item CORRECT, nor prove the recurrence has the same root cause as the original (a same-symptom recurrence with a genuinely different cause is a real possibility — §11.4.102 systematic-debugging decides that, and the linked-

and-reopened item is where that investigation is recorded, not a reason to fork a new id). This anchor BINDS rather than re-authors: the intake seam is §11.4.202's, the keying rule and negative-control discipline are §11.4.186's, the reopen mechanics and attribution are §11.4.34's, the counter and its history document are §11.4.55's, id stability is §11.4.54's, and the priority consumers are §11.4.132/ §11.4.189.

Classification: universal (§11.4.17) — references no project-specific hardware, vendor, or tracker; the consuming project supplies its tracker, its id scheme, its normalisation function, and its intake wiring per §11.4.35. Composes §11.4.1 / §11.4.6 / §11.4.34 / §11.4.54 / §11.4.55 / §11.4.66 / §11.4.90 / §11.4.93 / §11.4.101 / §11.4.102 / §11.4.105 / §11.4.107(10) / §11.4.132 / §11.4.148 / §11.4.186 / §11.4.189 / §11.4.202. Propagation gate `CM-COVENANT-114-214-PROPAGATION` (literal `11.4.214`) + recommended gate `CM-RECURRENCE-LINKS-NOT-MINTS` (every intake path runs the dedup check before minting; the oracle's golden-good/ golden-bad/negative-control are wired into the meta-test; every CONFIRMED-duplicate link resolves to a canonical CHAIN HEAD, and every link whose resolved head was in a TERMINAL state at link time has a corresponding reopen event ON THAT HEAD). **The gate's own scope is load-bearing and is itself §11.4.201-bound:** it MUST NOT fire on a link whose resolved head was OPEN at link time (LINK-only is the CORRECT action there per the open-original case), nor on a clause-(3) UNDECIDED `candidate-duplicate` link (which is by design un-reopened, and is the mandated autonomous default), nor demand a reopen on an intermediate `duplicate-of` closure (clause (0) FORBIDS that reopen — a gate demanding it would enforce the very defect the anchor prevents) — a gate firing on any of these would be exactly the §11.4.201(1) false-positive refusal that halts correct work and teaches operators to bypass it, and it would condemn behaviour THIS anchor mandates. Its golden-FALSE fixture set MUST therefore include an open-head link, a candidate-duplicate link, AND a link whose immediate target is a `duplicate-of` closure with the reopen correctly landing on the chain head — the gate MUST NOT fire on any of the three (§11.4.201(3) / §11.4.107(10)); its golden-TRUE set MUST include a terminal-head link with NO reopen event (the gate MUST fire) (file a known recurrence and let it mint a fresh id without a link → the gate FAILs; strip the negative control so two distinct same-subject items merge → the oracle self-check FAILs; strip the literal → the propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md §11.4.214.

Non-compliance is a release blocker regardless of context. No escape hatch — no

```
--skip-dedup-at-intake , --mint-every-report , --bare-substring-  
dedup , --merge-without-negative-control , --recurrence-as-new-id-OK ,  
--reopen-counter-may-undercount flag exists.
```

2026 07 17

§ . . — A doc that BINDS work MUST live, tracked, in the repository where that work happens (research-derived, — —)

Forensic anchor (genericised, 2026-07-17). In a consuming project, the document that DEFINED a binding acceptance gate — the gate a body of work was ordered to satisfy before proceeding — was **untracked** and lived in an out-of-tree scratch directory belonging to a different working checkout, i.e. it existed in NO repository and in no other agent's or operator's checkout. The gate could therefore not be honoured (agents could not read it), not be reviewed (no diff, no history, no authorship), not be version-controlled (no revision, no rollback, no §11.4.44 header), and not be reconciled when it contradicted the source of truth. A binding artifact that lives nowhere is not governance — it is a rumour with an imperative mood, and work ordered against it cannot be verified against it by anyone but the one process that happened to hold the file.

The mandate. Any document that BINDS work — one that defines a gate or acceptance criterion, states a contract the work is judged against, or carries a plan/order the work is required to satisfy — MUST be tracked, in the repository where that work happens, at a declared path, BEFORE it may be cited as binding. **(1) BINDING ⇒ TRACKED, and un-tracked ⇒ NOT BINDING.** A binding claim sourced from an untracked or out-of-tree document MUST NOT be enforced against any work; the correct response to "the gate is defined in a file nobody can read" is to land the document (or decline the gate), never to obey a citation of it. **(2) IN THE REPOSITORY THE WORK HAPPENS IN.** Tracked in a sibling checkout, a scratch directory, a chat message, an issue comment, or a personal note is NOT tracked for this purpose — the test is whether an agent or operator who clones ONLY the repo the work happens in can read the binding text at its declared path. Where a binding doc must govern several repositories, it is inherited BY REFERENCE from a shared, tracked, version-controlled dependency (§11.4.28 / §11.4.177), never copied and never left out-of-tree. **(3) THIS IS THE PRIOR**

QUESTION §11.4.212 CANNOT ASK. §11.4.212 mandates that every doc in the export scope be README-reachable — but its check ENUMERATES the in-repo scope and asserts each member is reachable, so a document that exists in no checkout is never enumerated and the gate cannot FAIL on it. §11.4.212 answers "is it findable?"; §11.4.215 answers "is it here at all?". Once a binding doc is in the repo, §11.4.212 (reachable from the README), §11.4.44 (revision header), §11.4.65/§11.4.73 (exports), and §11.4.106 (kept in sync) all attach to it as normal — this anchor is their precondition, not their competitor. **(4) THE**

GENERALISATION OF §11.4.95. §11.4.95 is this exact rule scoped to ONE named artifact: an authoritative tracker database is TRACKED, NEVER gitignored, "NOT a build artefact, it IS authoritative source data", as an explicit named carve-out from the gitignore rule. §11.4.215 widens the artifact class from that one file to **any doc that binds work**, exactly as §11.4.212 was minted as the strict generalisation of §11.4.57. §11.4.11 is the counterweight and stays intact: logs, forensic captures, and scratch artifacts remain out-of-tree and untracked by default — §11.4.215 does NOT track them, and a doc's being untracked is only a defect once it is CITED as binding. That is the discriminator: **binding is what pulls a document into the repository**, and a document nobody cites as binding is governed by §11.4.11, not by this anchor.

Honest boundary (§11.4.6). In-repo presence proves the binding text is AVAILABLE, ATTRIBUTABLE, and VERSION-CONTROLLED — it proves nothing about whether the text is CORRECT, current, or self-consistent (those remain §11.4.44 / §11.4.106 / §11.4.186), and nothing about whether the gate it defines is well-seated (that is §11.4.120's seam-placement clause, whose forensic case was defined by this same out-of-tree document — the two defects were carried by one artifact, which is itself the argument that an unreviewable binding doc is where bad gates hide).

Classification: universal (§11.4.17) — references no project-specific hardware, vendor, or layout; the consuming project supplies its repository, its declared paths for binding docs, and its shared-dependency mechanism per §11.4.35. Composes §11.4.11 / §11.4.28 / §11.4.30 / §11.4.35 / §11.4.44 / §11.4.57 / §11.4.65 / §11.4.73 / §11.4.95 / §11.4.106 / §11.4.120 / §11.4.177 / §11.4.186 / §11.4.197 / §11.4.212.

Propagation gate `CM-COVENANT-114-215-PROPAGATION` (literal `11.4.215`) + recommended gate `CM-BINDING-DOC-IN-REPO` (every doc cited as binding by a gate, order, or acceptance contract resolves to a tracked path in the repository the work happens in — or is inherited by reference from a tracked shared dependency;

a binding citation resolving to an untracked or out-of-tree path FAILs) + paired §1.1 mutation (move a binding doc out of the tree, or untrack it, while a gate still cites it → the gate FAILs; strip the literal → the propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.215.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--binding-doc-out-of-tree` , `--untracked-gate-definition` , `--cite-doc-nobody-can-read` , `--scratch-dir-is-good-enough` , `--binding-by-rumour` flag exists.

§26 07 22 . . — Canonical machine-readable design-token source: one CSS custom-property file (`:root` + `[data-theme="dark"]`), generated bindings everywhere else (User mandate, — —)

Forensic anchor — proving project (Helix Thready MVP design package, 2026-07-22). The mandate's shape was proven end-to-end in the Helix Thready design package (`docs/public/research/mvp/design/`): ONE machine-readable token file (`opendesign/tokens.css`) carries every color / type / spacing / radius / elevation / motion value as CSS custom properties — light values in a `:root {}` block, dark values in a `[data-theme="dark"]` override — and every value is evidence-traceable, never invented: the brand chartreuse `#B6E376` is the recorded eyedrop MEAN of 1,625,855 green-dominant logo pixels, the mint secondary `#ABDDC9` the mean of 618,886 mint-region pixels, the light accent MEASURED at 6.03:1 on white and the dark accent at 13.56:1 on the night surface. Every non-CSS platform in the fan-out consumes a GENERATED binding (token → Compose `Color / Dp` , Lipgloss `lipgloss.Color` , Flutter/Qt theme, Figma Variables payload), never a hand-kept copy; the not-yet-built codegen bridges are honestly `[OPEN]` -tracked, never claimed shipped.

The mandate. Every project governed by this Constitution that ships user-facing interfaces MUST maintain its design tokens as ONE canonical, machine-readable source, and everything else MUST derive from it: (1) **Canonical file shape** — CSS custom properties, light values in `:root {}` , dark values via the `[data-theme="dark"]` override pattern, composing the project's sanctioned dark-mode

mechanisms (`@media (prefers-color-scheme: dark)` system default, explicit `data-theme` attribute, framework-compact class) — never duplicate hand-kept light/dark palettes. (2) **Generated bindings only** — every consuming surface (web CSS, native mobile, desktop, TUI, design tools, rendered doc artifacts) binds to the canonical source either directly via `var(--...)` or through a GENERATED binding produced by codegen FROM the canonical file; a hand-copied or hand-restated token value anywhere (an inline hex that restates a token, a hand-kept per-platform palette) is a §11.4.216 violation — hand-copies drift silently, and the drift is invisible until a surface renders off-brand. (3) **No invented values — §11.4.6 applied to design** — every color / size / duration in the token source is traceable to captured evidence: brand colors to a recorded eyedrop / measurement of the canonical brand asset, accent pairs to a MEASURED WCAG contrast ratio against their PAIRED background, structural values to the documented shared-system scale; a hex that appears in any design artifact and resolves to no token is a violation — flag it and use the nearest existing token (per §11.4.162's no-invented-values discipline), never mint an ad-hoc value. (4) **The token source is the white-label seam** — per-brand / per-account re-theming swaps token VALUES within the §11.4.217 closed override set, never component structure.

Honest boundary (§11.4.6). One canonical source proves consistency of DEFINITION, not correctness of USE — a component can still misuse a token (brand-as-body-text); role discipline is §11.4.217's and rendered-pixel proof is §11.4.170's. And "generated binding" follows §11.4.197 wiring honesty: a codegen bridge that does not exist yet is an `[OPEN]` -registered gap (§11.4.223), never claimed working.

Classification: universal (§11.4.17) — references no project-specific brand or platform; the consuming project supplies its token-file path, brand-asset evidence, platform fan-out list, and codegen bridges per §11.4.35. Composes §11.4.6 / §11.4.162 (OpenDesign is the mandated token SYSTEM — this anchor pins its canonical FILE shape and consumption discipline) / §11.4.170 / §11.4.190 / §11.4.197 / §11.4.217 / §11.4.223. Propagation gate `CM-COVENANT-114-216-PROPAGATION` (literal `11.4.216`) + recommended gate `CM-DESIGN-TOKEN-SINGLE-SOURCE` (the canonical token file exists with `:root {} + [data-theme="dark"]` blocks; tracked design artifacts bind via `var()` rather than restating token hex inline; every declared platform names its generated-binding

artifact or its [OPEN] id) + paired §1.1 mutation (hand-estate a token hex in a tracked artifact → gate FAILs; strip the dark override block → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.216.

Non-compliance is a release blocker regardless of context. No escape hatch — no
`--hand-copied-palette-OK`, `--inline-hex-OK`, `--invent-token-value`, `--light-only-tokens`, `--platform-keeps-own-palette` flag exists.

§ . . — OpenDesign brand contract:
every UI-shipping product carries a -section
`DESIGN.md` + `tokens.css` twin with disjoint color
roles, AA-pinned accents, and a locked
attribution footer (User mandate,
- -)

Forensic anchor — proving project (Helix Thready MVP design package, 2026-07-22). Helix Thready ships `design/opendesign/DESIGN.md` in the OpenDesign 9-section brand-contract schema — VERIFIED against the tool's own source per §11.4.99 (spec, `design-systems/README.md`, the contributor authoring guide, and a measured 150-file corpus: two coexisting heading families, both parsed by number prefix only) — plus its machine-readable twin `tokens.css` per the OpenDesign token schema, and an honest `TOOLING.md` capability check. The contract encodes three DISJOINT color roles (brand decorative, accent interactive, semantic state), the closed white-label override set the three shipped sibling themes prove works, and the locked attribution footer.

The mandate. Every product governed by this Constitution that ships user-facing interfaces MUST carry an OpenDesign brand contract under its design area (`design/opendesign/` or the project's declared equivalent per §11.4.35), consisting of: (1) **DESIGN.md in the verified 9-section schema** — H1 product title with the `> Category:` blockquote IMMEDIATELY after the H1 (the picker parses the line immediately after the H1), nine numbered sections (the parser matches the number prefix, not the heading text), a "Font labels for catalog extraction" block (`Display:` / `Body:` / `Mono:` lines) in the typography section, real hex codes, CSS variables wrapped in `:root {}` (never bare) with dark tokens via the `[data-theme="dark"]` override pattern, real component CSS with `prefers-`

reduced-motion targeting specific elements (never a global *), and :focus-visible on every interactive component — the tool's own blocking review lens, satisfied in-file. Format claims are verified against the tool's CURRENT source per §11.4.99, never assumed from memory. (2) **The tokens.css machine-readable twin** — the §11.4.216 canonical token file (or a compiled view of it), following the tool's token schema; prose and tokens never diverge. (3) **Disjoint color roles by contract** — brand is DECORATIVE only (never body text; the Thready brand base measures 1.47:1 on white — decorative by physics, not by preference), accent is INTERACTIVE and AA-PINNED (each accent/on-accent pair carries a measured $\geq 4.5:1$ ratio against its paired background, both modes), semantic (success / warn / danger) is STATE and is NEVER overridden by a brand — a brand color can never mask an error signal. (4) **White-label closed override set, server-validated** — per-account branding may swap decorative brand tokens, the accent pair, logo, and slogan; accent overrides are SERVER-VALIDATED to WCAG AA $\geq 4.5:1$ or REJECTED (the rejection carries the measured ratio + an AA-passing suggestion — instructive, never bare); structural tokens, neutral surfaces, and semantic tokens are NOT account-overridable (advanced neutral overrides gate behind a full AA re-validation). (5) **Locked attribution footer** — the organization's attribution footer ships on every surface and is NEVER removed, even fully white-labeled.

Honest boundary (§11.4.6). The contract governs DEFINITION and roles; it does not prove rendered correctness (§11.4.170 / §11.4.190) nor replace §11.4.162's token-system, test-coverage, and extend-don't-reimplement obligations — this anchor pins the contract ARTIFACT §11.4.162 consumes. Schema-version drift between the local tool checkout and the doc is recorded as fact and re-verified at integration, never papered over.

Classification: universal (§11.4.17) — the consuming project supplies its brand values, category line, white-label storage/API bindings, and attribution string per §11.4.35. Composes §11.4.6 / §11.4.99 / §11.4.162 / §11.4.190 / §11.4.216 / §11.4.220 / §11.4.223. Propagation gate CM-COVENANT-114-217-PROPAGATION (literal 11.4.217) + recommended gate CM-OPENDESIGN-BRAND-CONTRACT (DESIGN.md present with the 9 numbered sections + category blockquote + font-labels block + :root {} + [data-theme="dark"] + targeted reduced-motion; tokens.css twin present and consistent; accent pairs carry measured ratios; the attribution footer token/lockup is present and not account-overridable) + paired §1.1

mutation (drop the font-labels block, or let a brand token override `--danger`, or mark the footer overridable → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.217.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--skip-brand-contract`, `--prose-only-design-md`, `--brand-as-text-OK`, `--unvalidated-accent-override`, `--remove-attribution-footer`, `--semantic-tokens-brandable` flag exists.

§ 2026 07 — **Living design-library catalogue: every component × state × theme rendered self-contained + a cross-platform reusability matrix with per-cell provenance (User mandate, — —)**

Forensic anchor — proving project (Helix Thready MVP design package, 2026-07-22). Helix Thready's `design/library/` renders the ENTIRE component set as a living catalogue: `components.html` — 14 groups, ~38 components, ≈ 118 component-state cells, each auditable in BOTH themes via a real toggle (≈ 236 rendered states), zero external requests, interactive behaviors REAL not mocked (theme cycling stamps `data-theme` and persists; column sorting drives `aria-sort`; a real `showModal()` dialog; an auto-dismissing toast that pauses on hover) — plus `components-sheet.svg` (an importable one-sheet) and `platform-map.md`, a 36-components × 8-platforms reusability matrix where EVERY cell carries a VERIFIED / PROPOSED / ASSUMED provenance marker re-checked at source.

The mandate. Every project governed by this Constitution that ships user-facing interfaces MUST maintain a LIVING design-library catalogue: (1) **Complete component × state × theme rendering** — every component of the project's design system, in every declared state (default / hover / focus-visible / disabled / loading / error / empty / skeleton as applicable), auditable in BOTH light and dark from one artifact; a component or state that exists in the spec but not in the catalogue is a coverage gap, tracked per §11.4.223. (2) **Self-contained artifacts** — the catalogue renders with ZERO external requests: no CDNs, no external fonts (self-hosted or system-stack fallback stated honestly), vendored libraries inlined; a catalogue that

needs the network is not auditable offline and violates the project's CSP/offline posture. (3) **Real behaviors, token-bound** — interactive exhibits are real (real dialog, real sort, real toggle), every color/size binds to the §11.4.216 token source, and all motion honors `prefers-reduced-motion` per §11.4.221. (4) **Cross-platform reusability matrix with per-cell provenance** — a components × platforms matrix where EVERY cell is marked from the closed provenance set: VERIFIED (reproduced from source and re-checked at source, with the check date), PROPOSED (a `[DEFAULT – adjustable]` design pending review), or ASSUMED (inference, labeled as such); an unmarked cell reads as ASSUMED. The matrix is read BEFORE implementing any component on any platform; scaffolds are named scaffolds — a scaffold cell claiming VERIFIED is a §11.4 PASS-bluff at the design-planning layer.

Honest boundary (§11.4.6). The catalogue proves the DESIGN renders and behaves as specified in a browser-class runtime; it does not prove any platform IMPLEMENTATION (that is the matrix's per-cell provenance + each platform's own §11.4.170 rendered-pixel proof), and PROPOSED components are named proposals pending the upstream-vs-project split decision — never silently promoted to shipped.

Classification: universal (§11.4.17) — the consuming project supplies its component set, state vocabulary, platform list, and matrix location per §11.4.35. Composes §11.4.6 / §11.4.162 / §11.4.170 / §11.4.216 / §11.4.221 / §11.4.223. Propagation gate `CM-COVENANT-114-218-PROPAGATION` (literal `11.4.218`) + recommended gate `CM-DESIGN-LIBRARY-LIVING-CATALOGUE` (the catalogue artifact exists, is self-contained — zero external URL references —, covers every spec'd component × state in both themes or tracks the gap, and the reusability matrix carries a closed-set provenance marker in every cell) + paired §1.1 mutation (introduce a CDN reference → gate FAILs; blank a matrix cell's marker or claim VERIFIED on a scaffold row → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.218.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--cdn-in-catalogue`, `--single-theme-catalogue`, `--mocked-behaviors-OK`, `--unmarked-matrix-cell`, `--scaffold-claims-verified` flag exists.

§ . . — Per-app screen catalogues: full-IA coverage per surface, per-platform chrome variants, honesty markers for unbuilt behavior (User mandate, — —)

Forensic anchor — proving project (Helix Thready MVP design package, 2026-07-22). Helix Thready's `design/screens/` carries a screen catalogue PER SURFACE — 13 web screens + an interactive prototype shell walking the four key journeys, 11 device-framed 390×844 mobile artifacts with an Android/iOS chrome TOGGLE (+ HarmonyOS/Aurora port notes) covering EVERY node of the wireframes' information architecture, a desktop shell with per-OS title bars / menus / tray / file-drop mocks, ten 80-column TUI mockups covering the full state machine, and a 3-page marketing set with a claims-traceability register. Undefined product behavior is rendered HONESTLY: placeholder billing amounts labelled as placeholders, an unspecified forgot-password flow linked with an inline note instead of a fake flow, an undecided channel source rendered as a DISABLED tile — each carrying its `[OPEN: id]` .

The mandate. Every application/surface a governed project ships MUST have a screen catalogue: (1) **Full-IA coverage** — the catalogue covers EVERY node of the surface's information architecture (every screen of the wireframe/IA contract); a screen present in the IA but absent from the catalogue is a coverage gap tracked per §11.4.223 — never silently skipped. (2) **Per-platform chrome variants** — where one design serves several host platforms, the catalogue renders the platform chrome variants (mobile OS chrome toggle, per-OS window chrome, terminal-width constraints) rather than pretending one chrome fits all; ports not yet designed are tracked, not faked. (3) **Self-contained, token-bound, both themes** — each screen is a self-contained artifact binding the §11.4.216 tokens, rendered light + dark, populated with REALISTIC product data (never lorem-ipsum for load-bearing content), carrying its per-screen state contract (loading / empty / error / populated states declared). (4) **Honesty markers for unbuilt behavior** — where the backing backend, contract, or product decision does not exist, the screen MUST say so: a `[GAP: id]` / `[OPEN: id]` marker (registered per §11.4.223), a labelled placeholder, an inline note, or a disabled control — NEVER a designed-as-if-real fake flow; designing a flow for an undecided behavior and presenting it as settled is a §11.4 PASS-bluff at the product-design layer.

Honest boundary (§11.4.6). A screen catalogue is the DESIGN contract, not the implementation: it proves every IA node has a designed, state-complete, theme-complete rendering — implementation correctness stays with §11.4.170 host-rendered pixel proof, §11.4.190 web-quality proof, and the §11.4.185 manual-QA final gate.

Classification: universal (§11.4.17) — the consuming project supplies its surface list, IA contract, chrome-variant matrix, and realistic-data corpus per §11.4.35. Composes §11.4.6 / §11.4.162 / §11.4.170 / §11.4.185 / §11.4.190 / §11.4.216 / §11.4.218 / §11.4.223. Propagation gate `CM-COVENANT-114-219-PROPAGATION` (literal `11.4.219`) + recommended gate `CM-SCREEN-CATALOGUE-FULL-IA` (every IA node maps to a catalogue screen or a registered `[OPEN]` id; every screen renders both themes; every screen whose backing contract is absent carries its honesty marker) + paired §1.1 mutation (remove a screen without registering the gap → gate FAILs; strip an honesty marker from an unbacked flow → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.219.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--partial-ia-coverage`, `--one-chrome-fits-all`, `--fake-flow-OK`, `--lorem-ipsu`^{11 4 220}`m-content`, `--single-theme-screens` flag exists.

§ . . — Open-first design tooling: self-hosted open platforms are primary; proprietary tools are import/export targets only; design sources live in-repo in open formats (User mandate, — —)

Forensic anchor — proving project (Helix Thready MVP design package, 2026-07-22). Helix Thready's tooling capability check (`opendesign/TOOLING.md`) and export capstone prove the posture: every design source lives IN-REPO in open formats (7 SVG brand masters, `tokens.css`, JSON manifests + 28 genuine `.od-figma.json` IR captures, 6 hand-authored Lottie JSONs, self-contained HTML screens); the OpenDesign daemon/CLI status was RUNTIME-PROBED (live health endpoint, CLI round-trip, headless-vs-operator matrix), never assumed; and Figma is an IMPORT-ONLY bridge — the import plugin pins `networkAccess: ["none"]`,

no cloud token exists, and the docs state verbatim that any "pushed to Figma cloud" claim from that environment would be a fabrication. The actual proprietary-tool file is `[OPEN]` -tracked, not claimed.

The mandate. Every governed project's design workflow MUST be OPEN-FIRST:

(1) **Self-hosted open platform primary** — the collaborative design platform is an open-source, self-hostable tool (PenPot as the standing default), run SELF-HOSTED; together with the §11.4.162 OpenDesign engine it forms the primary design toolchain. (2) **Consumption ONLY via the PenPot submodule** — ALL PenPot mechanisms and utilities (the deployment recipe, RPC client utilities, access-token minting, import tooling, browser-hydration automation, MCP wiring) live in the dedicated reusable PenPot submodule (working name `vasic-digital/PenPot` ; the final repository name was UNCONFIRMED at this anchor's writing — the generic referent "the PenPot submodule" is normative until the name lands), inherited BY REFERENCE per §11.4.28; the PenPot submodule itself consumes the Containers submodule (§11.4.76, rootless per §11.4.161) for orchestration and DECLARES that dependency in its §11.4.31 `helix-deps.yaml` manifest. Projects MUST NOT hand-roll PenPot deployment or client code outside the submodule — a missing capability is extended UPSTREAM in the submodule per §11.4.74, never worked around locally. The submodule carries MECHANISMS/UTILITIES ONLY, never project design content — per-project design sources stay in the project repo per clause (4). (3) **Proprietary tools are optional bridges, never the source of truth** — proprietary design tools (Figma or equivalent) MAY be used as import/export/review targets via documented bridge kits (variables payloads, IR captures, import guides), but NO design value may exist ONLY inside a proprietary tool or its cloud: the moment a decision lands, its canonical form lands in-repo. (4) **In-repo open formats** — design sources are versioned in the repository in open, diffable formats: SVG (vector masters), CSS (tokens per §11.4.216), JSON (manifests, IR captures, variables payloads), Lottie JSON (motion per §11.4.221), self-contained HTML (screens/library per §11.4.218/§11.4.219); binary proprietary formats are derived artifacts at most, never sources — the §11.4.215 discipline applied to design: a design that binds work MUST live tracked in the repo. (5) **Runtime-probed capability claims** — every tooling claim ("the daemon runs", "the CLI exports", "import works") is established by a runtime probe recorded as evidence (§11.4.201 real-condition discipline; §11.4.99 latest-source verification for format contracts), with an explicit headless-vs-operator-interactive matrix and an HONEST-LIMITS section; a cloud push claimed without a real credentialed run is a

fabrication of §11.4 PASS-bluff severity. (6) **PenPot automation operational rules (research-verified)** — (a) the HEADLESS write path is the token-authenticated **RPC API**: access tokens (minted under the backend's `enable-access-tokens` flag, handled per §11.4.10) are the MANDATED auth mechanism for all automation — interactive OAuth relays / session-cookie scraping are FORBIDDEN as automation auth; (b) API-created PenPot files REQUIRE a **browser hydration pass** before handoff — thumbnails and text layout are materialized by the browser runtime, not by the write API, so a handoff of an unhydrated API-created file claimed render-ready is a §11.4 PASS-bluff at the design-handoff layer; the hydration pass is automated via the PenPot submodule's browser-hydration utility and its execution is recorded as evidence (§11.4.69).

Honest boundary (§11.4.6). Open-first governs WHERE truth lives and WHICH tool is primary; it does not forbid operators using proprietary tools for review/hand-off, and it does not claim bridges exist before they are built — missing bridges (PenPot/Lottie export, proprietary-file creation) are `[OPEN]` -registered work per §11.4.223/§11.4.197; and until the PenPot submodule's final repository name is confirmed, citations use the generic referent honestly rather than asserting an unconfirmed name as fact.

Classification: universal (§11.4.17) — the consuming project supplies its platform choice within the open-first constraint, its PenPot-submodule pointer + configuration (instance URL, token env-var NAMES per §11.4.10), bridge kits, and probe scripts per §11.4.35. Composes §11.4.6 / §11.4.10 / §11.4.28 / §11.4.31 / §11.4.69 / §11.4.74 / §11.4.76 / §11.4.99 / §11.4.161 / §11.4.162 / §11.4.201 / §11.4.215 / §11.4.216 / §11.4.223. Propagation gate `CM-COVENANT-114-220-PROPAGATION` (literal `11.4.220`) + recommended gate `CM-OPEN-DESIGN-TOOLING-PRIMARY` (the self-hosted open platform is declared and consumed via the PenPot submodule — no hand-rolled PenPot deployment/client code in the project tree; the submodule's `helix-deps.yaml` declares the Containers dependency; automation authenticates via access tokens, never an OAuth relay; every API-created file handed off carries a recorded hydration-pass run; every design source referenced by the design area resolves to an in-repo open-format file; every tooling capability claim cites a recorded probe; no design artifact exists only as a proprietary-cloud reference) + paired §1.1 mutation (hand-roll a PenPot RPC client in a project tree, or declare a proprietary tool as the source of truth, or hand off an API-created file

with no hydration-pass record → gate FAILs; claim a capability with no recorded probe → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.220.

Non-compliance is a release blocker regardless of context. No escape hatch — no

`--figma-is-source-of-truth` , `--cloud-only-design` , `--proprietary-format-source` , `--hand-rolled-penpot-stack` , `--oauth-relay-automation` ,
`--skip-hydration-pass` , `--unprobed-tooling-claim` , `--rootful-design-stack` flag exists.

2026 07 22

§ . . — **Motion discipline: token-bound durations/easings + machine-readable manifest + reduced-motion static fallbacks everywhere** (User mandate, — —)

Forensic anchor — proving project (Helix Thready MVP design package,

2026-07-22). Helix Thready's `design/motion/` package proves the discipline: six hand-authored vector Lottie animations bound to the token motion values, each validated by parse + structural schema check + REAL render (lottie-web under jsdom with animation-progression assertions, plus headless-Chromium light+dark contact sheets) — unverified runtimes (`lottie-compose` / `lottie-ios`) honestly `[OPEN]` -tracked; FIVE static reduced-motion fallback SVGs whose geometry is DERIVED FROM THE LOTTIE KEYFRAMES (the exact poster frame for loopers, the final frame for one-shots); and a machine-readable `motion-manifest.json` v2 naming, per animation, its loop mode, poster frame, segment markers, and `reducedMotionFallback` — with one `DELIBERATE null` (an instant-cut transition) recorded with its reason, never a silent absence.

The mandate. All motion in every governed project MUST satisfy: (1) **Token-bound timing** — every duration and easing binds to the §11.4.216 motion tokens (the shared base durations + easing curves); documented choreography MAY extend them, but an animation timed by an ad-hoc literal that bypasses the token layer is a §11.4.221 violation (per §11.4.6 no-invented-values). (2) **Reduced-motion fallback per animation** — every playable animation ships a STATIC fallback derived from its own keyframes (poster frame for looping animations, final frame for one-shots), delivered as an open-format asset; where the correct

reduced-motion behavior is intentionally something else (instant cut, no artwork), that is recorded EXPLICITLY in the manifest with its reason — a missing fallback is a gap, a silent missing fallback is a violation. (3) **Machine-readable manifest** — a versioned motion manifest enumerates every animation with its id, loop mode, poster frame, reduced-motion mode, and fallback reference; runtimes consume the MANIFEST, not folklore. (4) **prefers-reduced-motion honored everywhere** — every surface (web, library, screens, previews, implementations) gates motion on `prefers-reduced-motion: reduce` with TARGETED component-level gates (never a global * kill-switch, per the §11.4.217 review lens), swapping in the manifest-named fallback. (5) **Render-validated, honestly bounded** — no animation is claimed working without a real render in a real runtime (the §11.4.107(10) analyzer discipline applied to motion assets: parse + structure + rendered geometry with progression assertions); runtimes not yet exercised are named `[OPEN]` items, never assumed.

Honest boundary (§11.4.6). The manifest + fallbacks prove reduced-motion COVERAGE; they do not prove physical-display fidelity (a pixel A/B of statics against frozen frames on a real display remains its own tracked item), and jsdom/headless renders do not certify mobile runtimes — the boundary is stated, tracked, and never blurred.

Classification: universal (§11.4.17) — the consuming project supplies its animation inventory, motion tokens, manifest path, and runtime list per §11.4.35. Composes §11.4.6 / §11.4.107(10) / §11.4.162 / §11.4.170 / §11.4.216 / §11.4.218 / §11.4.223. Propagation gate `CM-COVENANT-114-221-PROPAGATION` (literal `11.4.221`) + recommended gate `CM-MOTION-TOKEN-BOUND-REDUCED-FALLBACK` (every animation's timing resolves to tokens; the manifest enumerates every playable animation; every manifest entry names a fallback or an explicit recorded null-with-reason; every animation carries its parse/structure/render validation record) + paired §1.1 mutation (add an animation with a literal ad-hoc duration → gate FAILs; drop a fallback reference without a recorded reason → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.221.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--ad-hoc-duration`, `--no-reduced-motion-fallback`, `--global-star-motion-kill`, `--manifest-optional`, `--unrendered-animation-claimed-working` flag exists.

§ 11.4.22.2 — **Design-completion export wave: a design is not done until token codegen, the honest-output raster matrix, both-theme renders, hand-off bundles, and doc twins land (User mandate, — —)**

Forensic anchor — proving project (Helix Thready MVP design package, 2026-07-22). Helix Thready's export capstone proves the wave: an icon/asset raster matrix locked 1:1 to its generator (`generate-raster.sh`), which SELECTS its SVG renderer by PROBE — each candidate (rsvg-convert → inkscape → resvg → headless Chromium → ImageMagick) must pass a real probe render with content verification before it is used, a failing candidate falls through with a note, and if NONE passes the script exits with instructions and writes NOTHING ("No PNGs were written (output is not faked)"); an absent bundler (iconutil for `.icons`) SKIPS with a note — a verified live run produced 83 REAL rasters, every one checked non-blank; 49 screen PNGs (22 screens × light+dark at a confirmed 2× scale, theme forced deterministically, every file verified `size-floor + stddev > 0`); a 54-page bound design book; and PenPot + Figma hand-off bundles whose impossible/operator-only legs are recorded as DEFERRED with reasons — never faked.

The mandate. A design (a design system, a brand, a screen set, or a design change wave) is NOT "done" until its EXPORT WAVE has run and its outputs are verified: (1) **Per-platform token codegen** — the §11.4.216 generated bindings are (re)generated from the canonical token source, or their absence is

[OPEN] -tracked. (2) **Full icon/asset raster matrix via an honest-output pipeline** — the per-platform pixel/format matrix (every OS surface) is documented, locked 1:1 to a generator script, and produced by a pipeline that (a) PROBE-TESTS every candidate renderer with a real content-verified probe render before trusting it (§11.4.201 — a renderer's presence is not its correctness; an unprobed renderer would silently emit blank/garbage icons), (b) falls through to the next candidate on probe failure, (c) SKIPS-with-note any step whose tool is genuinely

absent (§11.4.3 / §11.4.69 honest-skip discipline), and (d) NEVER writes a placeholder, blank, or faked output — every produced raster is verified non-degenerate (size floor + non-blank content check; multi-resolution bundles verified to contain their declared members, per §11.4.38's open-the-artifact discipline). (3) **Both-theme screen renders + bound design book** — every catalogue screen rendered in light AND dark at the declared scale with deterministic theme forcing, each render verified non-blank, and the set bound into the portable design book. (4) **Hand-off bundles** — the open-platform and proprietary-bridge bundles (§11.4.220) built with per-leg HONEST statuses: produced / deferred-operator-only / impossible-with-reason. (5) **Doc twins regenerated** — the design area's §11.4.65 / §11.4.73 export siblings are regenerated so no reader gets a divergent view. A design wave reported complete without its export wave — or with any faked/unverified export output — is a §11.4 PASS-bluff at the design-delivery layer.

Honest boundary (§11.4.6). The export wave proves the design's ARTIFACTS exist, are real, and are non-degenerate; it does not prove visual correctness on physical displays (§11.4.170 / §11.4.190 / §11.4.185 own that), and a wave rendered before later additions is honestly stamped with its coverage boundary — regeneration lands as a tracked downstream pass, never claimed retroactively.

Classification: universal (§11.4.17) — the consuming project supplies its raster matrix, generator, renderer candidates, book builder, and bundle targets per §11.4.35. Composes §11.4.3 / §11.4.6 / §11.4.38 / §11.4.65 / §11.4.69 / §11.4.73 / §11.4.201 / §11.4.216 / §11.4.219 / §11.4.220. Propagation gate `CM-COVENANT-114-222-PROPAGATION` (literal `11.4.222`) + recommended gate `CM-DESIGN-EXPORT-WAVE-COMPLETE` (the raster generator probe-tests renderers and content-verifies outputs; the matrix doc and generator agree 1:1; every screen has both-theme verified renders or a tracked gap; every bundle leg carries an honest status; the doc twins are in sync) + paired §1.1 mutation (make a renderer skip its probe, or write a placeholder PNG on renderer absence → gate FAILs; report the wave complete with a stale doc twin → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.222.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--skip-export-wave` , `--unprobed-renderer` , `--placeholder-raster-OK` , `--single-theme-renders` , `--fake-bundle-status` , `--stale-doc-twins` flag exists.

2026 07 22

§ . . — Provenance-marker discipline for design documentation + central open-item registry in the area index (User mandate, — —)

Forensic anchor — proving project (Helix Thready MVP design package, 2026-07-22). Every statement in the Helix Thready design area carries an inline provenance marker from a closed vocabulary (`[VERIFIED]` , `[OPERATOR]` , `[RESEARCH]` , `[DEFAULT – adjustable]` , `[ASSUMED]` , `[BUILD-NEW]` , `[IN-HOUSE: module]` , `[CONSTITUTION §x]` , `[GAP: id]` , `[OPEN: id]`); the area index maintains the CANONICAL open-item registry — 41 registered ids across nine families, VERIFIED AGAINST THE TREE in both directions (every id in the tree exists in the registry, no registry id is absent from the tree), with per-id state (open / narrowed / closed, each closure citing the closing revision) — and a verified-vs-assumption ledger separating measured fact from proposed default from named scaffold.

The mandate. Every design document in a governed project MUST carry provenance discipline: (1) **Mandatory marker vocabulary** — a closed, project-declared marker set whose minimum core is `[VERIFIED]` (established from source/runtime/measurement, with the evidence or check date), `[OPERATOR]` (an operator decision), `[DEFAULT – adjustable]` (a proposed default pending review), and `[OPEN: id]` (an unresolved item with a stable id); projects extend the set (`[RESEARCH]` , `[ASSUMED]` , `[BUILD-NEW]` , `[GAP: id]` , `[IN-HOUSE: ...]` , `[CONSTITUTION §x]`) per §11.4.35. A load-bearing claim with no marker reads as UNVERIFIED — and presenting an assumption as verified is a §11.4.6 violation of §11.4 PASS-bluff severity at the design-documentation layer. (2) **Every unresolved decision is an `[OPEN: id]`** — with a stable, never-renumbered id (§11.4.54 discipline), owned by a named file, and never papered over with an invented answer. (3) **Central registry in the area index** — the design area's index maintains the SINGLE canonical open-item registry: every `[OPEN: id]` in the tree appears in it, every registry row names its owning file(s) and state, and the registry

is TREE-VERIFIED in both directions on every consolidation pass (an id in the tree missing from the registry, or a registry ghost with no tree presence, is a violation — the §11.4.186 consistency discipline applied to design opens); closures cite the closing revision, and narrowed items record exactly what closed and what remains. (4) **Opens map to tracked work** — registry items map to the project's workable-item tracking (§11.4.15 / §11.4.54 / §11.4.93) so the §11.4.197 completion mandate applies: a design open is driven to closed-with-evidence or explicitly-closed-with-reason, never silently abandoned.

Honest boundary (§11.4.6). Markers prove the EPISTEMIC STATUS of each claim, not its truth — a [VERIFIED] marker is only as good as its cited evidence (§11.4.123), and the registry proves open-item ACCOUNTING, not progress (§11.4.197 owns completion). This anchor binds the DESIGN area to the same honesty regime the tracker anchors give code — it cites §11.4.15/.54/.93/.186/.197 rather than duplicating them.

Classification: universal (§11.4.17) — the consuming project supplies its marker-set extensions, id families, area-index path, and tracker binding per §11.4.35.

Composes §11.4.6 / §11.4.15 / §11.4.44 / §11.4.54 / §11.4.61 / §11.4.93 / §11.4.123 / §11.4.186 / §11.4.197 / §11.4.212. Propagation gate CM-COVENANT-114-223-PROPAGATION (literal 11.4.223) + recommended gate CM-DESIGN-DOC-PROVENANCE-MARKERS (every design doc declares/uses the marker vocabulary; the area index registry exists; a tree-grep of open-item ids equals the registry set in both directions; every registry row names owner + state) + paired §1.1 mutation (add an [OPEN: id] to a doc without registering it → gate FAILs; delete a registered id from the tree while the registry still lists it open → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md §11.4.223.

Non-compliance is a release blocker regardless of context. No escape hatch — no --unmarked-claims-OK , --assumption-as-verified , --unregistered-open-item , --registry-optional , --silently-abandon-open flag exists.

§ . . – Test-first (TDD) for ALL work, not only fixes + a minimum code-coverage floor that is NECESSARY-never-sufficient (User mandate, – –)

Forensic anchor — verbatim user mandate (2026-07-22):

"Any work we do MUST START by writing the test! We MUST FULLY comply and follow the TDD - The test driven development! It MUST BE applied for any change, implementation, wiring, tests, challenges and any other form of codebase we write! Bash scripts as well! Coverage with TDD MUST BE close to 100%, not less than 85% !!!"

The gap this closes. §11.4.43 mandates the test-first (RED → ... → GREEN → VERIFY → DOCUMENT) workflow but its operative text binds **fixes** only ("Every fix MUST follow the 5-step TDD-fix workflow"); §11.4.115 refines that RED step onto the broken artifact; §11.4.146 binds reproduce-first. NONE of them binds a brand-new feature, a new wiring/integration, a new gate, a new Challenge, or a NEW BASH SCRIPT — none of which is a "fix", so none of which inherits a test-first obligation today. Separately, §11.4.27 and §11.4.169 mandate 100% **test-TYPE** coverage (breadth: WHICH KINDS of test exist per feature × platform) — an orthogonal axis that states no **code-coverage** figure anywhere. §11.4.224 closes exactly those two holes and claims nothing else as new.

(A) TEST-FIRST BINDS ALL WORK, not only fixes. For EVERY work product with an executable surface — a change, a new feature, an implementation, a wiring/integration, a refactor, a pre-build gate, a §1.1 mutation, a Challenge, test-harness code itself, AND **bash / shell scripts** — the test is WRITTEN FIRST, RUN FIRST, and OBSERVED TO FAIL for the right reason before the implementation exists; only then is the implementation written to make it pass. Where a broken artifact already exists the RED baseline is taken ON THAT ARTIFACT per §11.4.115 (and its polarity switch converts the same source into the permanent §11.4.135 guard); where the work is genuinely new, the RED is the test failing against the absent/unimplemented behaviour, captured before the implementation lands. A test authored AFTER its implementation demonstrates only that the test AGREES with the code that already exists — never that it CATCHES the defect (the §11.4.43 PASS-bluff, here generalised to every work class). **What counts as**

"the test" per work class (§11.4.6 — named, never left to interpretation):

behaviour code → an executing test asserting the behaviour; a bash/shell script → an executing test that invokes the script through its REAL invocation path and asserts its real observable effect (exit status + emitted output + state delta), never a `bash -n` parse-check alone (parse-cleanliness is §11.4.67's separate obligation and proves nothing about behaviour); a pre-build GATE → its paired §1.1 mutation IS the test-first artifact, authored and observed to make the gate FAIL BEFORE the gate is trusted (a gate TRUSTED before its mutation was observed to FAIL is unvalidated instrumentation per §11.4.115(F) — the condition is observation-before-trust, NOT authorship order, so a gate and its mutation authored in the same commit are compliant when the mutation is observed to make the gate FAIL before the gate is relied upon); a Challenge → its scored assertion against captured evidence per §11.4.5/§11.4.69; a behaviour-bearing config/data change (an init-script line, a package-set manifest, an action-registry row — any committed DATA an executing consumer reads to decide behaviour) → an executing test that drives that CONSUMER through the changed value and asserts the changed observable behaviour (the clause-(F) executable-surface test reaches the config THROUGH its consumer — a config change is never parked under the prose-exempt class, which is strictly pure prose / governance / documentation).

(B) CODE-COVERAGE FLOOR — ≥ 85% MINIMUM, ~100% TARGET. Every governed codebase carries a measured code-coverage figure that MUST NOT fall below **85%**, with **as close to 100% as the domain permits** as the standing target (composing §11.4.169's "100% coverage as the domain permits" phrasing on its own test-type axis). The floor applies to the executable codebase INCLUDING bash/shell scripts — a shell corpus exempted from measurement is exactly the gap the operator's "Bash scripts as well!" closes.

(B.1) SEVEN-TYPE BREADTH — the 100% AIM across a canonical test-type enumeration (extension, 2026-08-15). The 85% floor is a LINE-COVERAGE proxy across the executable codebase; the 100% AIM is the STANDING TARGET across the seven-canonical-test-type × executable-surface product per operator mandate 2026-08-15. The seven canonical types every consuming project MUST cover (where applicable) are: (1) UNIT, (2) INTEGRATION, (3) END-TO-END (E2E), (4) FULL AUTOMATION, (5) SECURITY, (6) PERFORMANCE AND BENCHMARK, (7) ANTI-BLUFF — each carrying its own §11.4.27 mandatory-breadth obligation (which this anchor cross-references, and §11.4.27's clause (B) is amended in step). A project measuring 100% line coverage on UNIT tests alone

while omitting one of the seven applicable types is a §11.4.224(B.1) violation regardless of the line-coverage number (§11.4.201(8) metric-validated-against-done: the correct end-state is line coverage AT target ACROSS all seven applicable types, never on one type alone). Honest boundary §11.4.6 / §11.4.17: a test type genuinely inapplicable to a project (e.g. a library with no runtime surface has no E2E user-journey to drive) is HONESTLY marked as §11.4.3 SKIP-with-reason at the type-breadth axis, never silently omitted; the universal DISCIPLINE is unchanged, the applicability of each type is the consumer's §11.4.35 DATA. CM-COVERAGE-FLOOR is EXTENDED — the measured numerator STILL counts only RED-capable tests, PLUS the gate FAILs on missing test-type-breadth coverage per the seven-type enumeration (a project measuring 90% line coverage but with zero ANTI-BLUFF tests → gate FAILs); a paired §1.1 mutation that strips one enumerated type from the type-breadth check → the gate FAILs.

(C) COVERAGE IS NECESSARY, NEVER SUFFICIENT — the anti-gaming clause (§11.4.201(8) / §11.4/§1.1). A coverage percentage is trivially gameable and MUST NEVER be reported as evidence a change is correct: **MEASURED FACT (in-session, 2026-07-22, this host):** a line-granular coverage instrument reported an `if ...; then ...; else ...; fi` line as COVERED while the `else` branch was NEVER taken — the metric read fully-covered on a line half of whose behaviour was untested; and a test that EXECUTES a line while ASSERTING NOTHING raises the same number by the same amount as a test that proves the behaviour. Therefore: (1) the REAL bar remains the §11.4 anti-bluff covenant + §1.1 — every test MUST catch its own negation, PROVEN by a paired mutation that makes it FAIL; a line executed by an assertion-free test is NOT covered in the sense that matters; (2) the floor is a MINIMUM ON A MEASURABLE PROXY, never a certificate of correctness — citing "≥ 85% covered" as proof a change works is a §11.4 PASS-bluff at the metric layer; (3) the floor NEVER stands alone — it composes with the §1.1 mutation pair, the §11.4.5/§11.4.69/§11.4.107 captured-evidence bar, and the §11.4.108 runtime signature, and it does NOT displace any of them.

(D) METRIC VALIDATION AGAINST THE DEFINITION OF DONE (§11.4.201(8)) — split verdict, stated honestly. §11.4.201(8) requires every gating metric to be evaluated at the correct end-state and REFUSED if the correct end-state does not move it to target. **Directional validation, performed 2026-07-22:** at the correct end-state — every behaviour test-driven, every test catching its own negation — the lines implementing that behaviour ARE executed by their own tests, so the

metric DOES move to target; the correct end-state does NOT refute the floor, and the floor is therefore VALID as a NECESSARY condition. **The converse is REFUTED** by the clause-(C) measured gameability fact: reaching the floor does NOT imply the end-state — hence necessary-never-sufficient, and hence clause (C) is load-bearing rather than decorative. **OWED, not claimed (§11.4.197)**: the EMPIRICAL calibration of the specific 85% threshold against a given project's REAL corpus — whether a genuinely correct end-state on THAT corpus lands at 85%, higher, or lower — has NOT been performed and MUST be tracked as a §11.4.197 work item by every consuming project that enforces the floor; this anchor asserts no project has already filed it (asserting an unverified tracker state would be the §11.4.6 violation this clause exists to prevent). **Interim precedence** (mirroring the precedent §11.4.201(8) itself set for the §11.4.135 ratchet): until that calibration lands the $\geq 85\%$ floor KEEPS GATING — an uncalibrated floor blocking a merge is bounded, visible and operator-resolvable, whereas disarming the only mechanism that surfaces untested code silently re-opens the hole (the §11.4.101 reversible-safe choice); the calibration is OWED WORK, never a licence to switch the floor off.

(E) HOW COVERAGE IS MEASURED — per language, honestly (§11.4.6 / §11.4.35). The consuming project supplies its per-language coverage instrument as DATA per §11.4.35; the UNIVERSAL rule is that an instrument MUST exist and MUST actually be RUN — **a floor nobody can measure is an unmeasurable gate and is itself a §11.4.201 bluff**, and a percentage nobody measured is a §11.4.6 fabrication. **Measured host reality (2026-07-22, the project instantiating this anchor)**: Go coverage is available built-in (`go test -coverprofile + go tool cover`); C/C++ via `gcov` ; and **for bash/shell NO dedicated coverage tool was installed** — `kcov` , `bashcov` , `bats` , `shellspec` and `shunit2` were ALL probed ABSENT (`shellcheck` IS present but is a LINTER, not a coverage instrument, and satisfies nothing here). **PROVEN-FEASIBLE bash mechanism (measured in-session, never asserted)**: trace-based line accounting needs NO additional tooling — `PS4='+COV:${BASH_SOURCE##*/}:${LINENO}:' bash -x <script>` emits the executed line numbers, and `executed-lines ÷ executable` (non-blank, non-comment) source-lines yields line coverage; a probe run correctly identified an uncovered function body as uncovered, and subshell bodies ARE traced. **HONEST LIMITS of that mechanism (measured, stated as fact)**: it is LINE coverage, NOT branch coverage — the clause-(C) same-line `if/else` counted as covered with one branch never taken; `set +x` regions and traps are not accounted. It is therefore a FLOOR-measuring instrument only; where branch-

level assurance is required the project MUST install a real coverage tool (e.g. `kcov`) or record the shortfall honestly. **Where NO instrument exists** for a language on a host, the honest responses are to INSTALL one, or to record an §11.4.3/§11.4.69 SKIP-with-reason plus a tracked §11.4.197 item — NEVER an invented percentage, and NEVER a silently-dropped floor. **EXCLUSION-LIST FENCE (amendment, remediation round of 2026-07-22 — the §11.4.135 exemption pattern applied to the coverage corpus)**. The corpus-scope exclusion list is NOT an unconstrained consumer knob: an exclusion is legal ONLY via a CHECKED-IN exclusion list, with EVERY entry justified from a CLOSED class set — `generated-code` (output of a committed generator per §11.4.77), `vendored-third-party` (executable code the project does not author), `non-shipping-fixtures-and-golden-assets` (test fixtures / golden files that never execute in the shipped product) — and enumerated as honest gaps in the coverage report, never silent (the exact discipline §11.4.135 applies to topology exemptions: legal ONLY via a checked-in map, enumerated as honest gaps — never silent). Excluding FIRST-PARTY executable code additionally REQUIRES a tracked §11.4.197 work item naming the plan to bring it into scope. An unlisted, unjustified, or un-tracked-first-party exclusion VOIDS the measured figure — a consumer that excludes most of its first-party executable code and truthfully measures $\geq 85\%$ on the remainder has satisfied the operator's floor in letter and voided it in substance; that silent-evasion channel is exactly what this fence closes.

(F) HONEST BOUNDARY / SCOPE (§11.4.6). "All work" means work products with an EXECUTABLE surface. A pure prose / governance / documentation edit has no executable surface and is NOT coverage-scoped — its discipline is §11.4.44 / §11.4.65 / §11.4.106 / §11.4.212 plus the §11.4.142 universal review; quoting a coverage percentage for a prose document would be meaningless, and demanding one would be a §11.4.201(1) false-positive refusal. BUT a governance change that SHIPS an executable artifact (a gate, a hook, a script, a mutation) is FULLY in scope for both (A) and (B). Test-first does NOT replace the §11.4.108 four-layer verification, the §11.4.40 full-suite retest, or the §11.4.185 manual-QA final gate — it PRECEDES them. And §11.4.224's CODE-coverage axis is DISTINCT from §11.4.27 / §11.4.169's 100% TEST-TYPE coverage (breadth of test kinds per feature $\times$ platform): both apply simultaneously, neither substitutes for the other, and satisfying one is never evidence for the other.

Classification: universal (§11.4.17) — references no project-specific hardware, vendor, or language; the consuming project supplies its per-language coverage instruments, its corpus scope, its CHECKED-IN clause-(E)-fenced exclusion list, and its threshold calibration per §11.4.35. Composes §11.4 / §11.4.1 / §11.4.5 / §11.4.6 / §11.4.27 / §11.4.35 / §11.4.43 (generalises its fix-scoped test-first to ALL work) / §11.4.50 / §11.4.66 (brownfield adoption is an operator question) / §11.4.67 / §11.4.69 / §11.4.77 / §11.4.85 / §11.4.107 / §11.4.108 / §11.4.110 / §11.4.115 / §11.4.122 / §11.4.135 / §11.4.142 / §11.4.146 / §11.4.169 (orthogonal test-TYPE axis) / §11.4.185 / §11.4.197 / §11.4.201(8) / §11.4.202 / §1.1. Propagation gate `CM-COVENANT-114-224-PROPAGATION` (literal `11.4.224`) + recommended gates `CM-TEST-FIRST-ALL-WORK` (every commit adding an executable artifact carries its test, and that test's failing-first observation is recorded — a new script/gate/feature landing with no test, or with a test whose RED was never captured, FAILs) + `CM-COVERAGE-FLOOR` (a measured coverage figure exists per declared language, was produced by a real instrument run rather than asserted, and is $\geq$ the declared floor; a declared language with NO instrument FAILs rather than silently passing; **the measured numerator counts ONLY lines covered by RED-CAPABLE tests** — every test contributing to the measured corpus MUST carry a recorded failing-first run (its clause-(A) RED record) OR a paired mutation observed to make it FAIL, with sampling-based mutation verification an acceptable stated implementation and a contributing test carrying ZERO assertions rejected by construction, so assertion-free execution cannot raise the measured figure and ONE real RED-captured test plus assertion-free padding to 85% FAILs this gate; AND the gate FAILs on any corpus exclusion that is absent from the checked-in list, not justified from the clause-(E) closed class set, or first-party without its tracked §11.4.197 item) + paired §1.1 mutations (land an executable artifact with no test → `CM-TEST-FIRST-ALL-WORK` FAILs; record a coverage percentage with no instrument run behind it, or drop a language's measurement entirely, or pad the numerator with an assertion-free test, or slip an unjustified exclusion into the corpus scope → `CM-COVERAGE-FLOOR` FAILs; strip the literal → the propagation gate FAILs). **Brownfield adoption is DELIBERATELY UNDEFINED by this anchor (§11.4.66):** when `CM-COVERAGE-FLOOR` first lands on a project whose legacy corpus measures below 85%, the gate-code work item MUST surface the adoption question to the OPERATOR per §11.4.66 BEFORE first enforcement — candidate options to present: immediate hard floor (day-one RED until the corpus is raised) / a ONE-TIME monotone-decrease ratchet per the §11.4.135 precedent / per-corpus phase-in (per language

or module as each enters scope) / floor-on-changed-code-only with a scheduled full-corpus deadline — recording the operator's answer as consumer DATA (§11.4.35); this anchor does NOT autonomously weaken the operator's unconditional "not less than 85% !!!" with an invented ratchet (a §11.4.122-class decision the operator owns). Gate-code = a separate work item — the CONTRACT is landed here, the gate CODE is NOT claimed shipped (§11.4.6).

Canonical authority: constitution submodule `Constitution.md` §11.4.224.

Non-compliance is a release blocker regardless of context. No escape hatch — no

```
--test-after-OK , --skip-red-first , --bash-exempt-from-tests , --  
coverage-not-applicable , --below-floor-OK ,  
1 1 4 2 2 5  
--coverage-proves-correct , --unmeasured-percentage-OK , --assertion-  
free-test-counts flag exists.
```

§ . . — Scheduler-quota burst-throttling telemetry + interactive-scope isolation from bursty fleets: progressive interactive degradation is diagnosed from throttling accounting over time, never from averages, and an interactive-latency-sensitive process never shares a fixed-quota no-burst scope with a bursty fleet (research-derived, - -)

Forensic anchor (genericised, 2026-07-22 — two independent instances in one day). (a) Progressive interactive sluggishness — worsening with the number of live subagents — was root-caused to a scheduler CPU quota (a cgroup `cpu.max` limit) with NO burst allowance on a session scope hosting BOTH the operator's interactive terminal-multiplexer server AND a spawning agent fleet. Measured live on a 64-CPU host: the scope's quota was 9.6 CPU per 100 ms enforcement period; average demand was 4.35 CPU — comfortably UNDER quota — yet 51 of 301 periods (16.9%) were throttled during active phases, ~6.4 runqueue-seconds of forced idle per wall-second at peak, 0% throttled during quiet phases (episodic, fleet-phase-correlated), and the lifetime ratio (10.9%) vs the live ratio (16.9%) showed it WORSENING with session age. The terminal server was co-resident in the throttled scope: it froze for ≥100 ms per throttled period, and consecutive throttled periods compounded into multi-second interactive stalls. Every

conventional instrument read HEALTHY — CPU average under quota, load average unremarkable, memory fine, thread counts fine — the degradation was INVISIBLE to averages and lived entirely in the throttling accounting (`cpu.stat` `nr_throttled` / `throttled_usec` `deltas`). (b) A sibling session launcher shipped a FIXED default quota (2 CPUs) on the same 64-core host — 18.8% of periods throttled and 1757 s of cumulative forced idle — fixed by computing the default from real detected host capacity (a per-core fraction with a floor and a cap), proven by a four-quadrant polarity test reading the LIVE cgroup back.

Two operative halves, BOTH mandatory:

(A) DIAGNOSTIC — throttling accounting over time under load, never averages. Any investigation of progressive / bursty / load-correlated interactive degradation (sluggish TUI or terminal, editor lag, stalls that worsen with parallel-agent or worker count) MUST sample the scheduling scope's THROTTLING ACCOUNTING — on Linux cgroup v2, `cpu.stat` `nr_throttled` + `throttled_usec` as DELTAS over a measured window under load, expressed as a fraction of elapsed enforcement periods, WITH a quiet-phase control sample (the §11.4.201(7)(b) needle in time: a mechanism that fires 17% under load and 0% quiet is proven load-coupled, not ambient) — BEFORE any verdict on the cause is minted. A HEALTHY AVERAGE WITH A HIGH THROTTLE RATIO IS THE TRAP: quota enforcement is per-period (canonically 100 ms), so a workload can sit far below its quota ON AVERAGE while a large fraction of individual periods hit the ceiling and stall every co-resident process; CPU averages, load average, memory, and thread counts are all structurally BLIND to it — their quiet "healthy" readings are §11.4.201(6) FALSE-NULLS, and "the averages look healthy, so the host is fine" without throttle-accounting deltas is a §11.4.6 guess. The same discipline applies to EVERY quota'd scheduler (container CPU limits, VM caps, hypervisor shares / steal time): the throttle/steal accounting IS the authoritative source (§11.4.201); the utilisation average is a proxy that the real condition does not have to satisfy.

(B) DESIGN — an interactive-latency-sensitive process NEVER shares a fixed-quota, no-burst scope with a bursty fleet. An interactive-latency-sensitive process (terminal-multiplexer server, TUI, editor, operator shell) MUST NOT be co-resident in a fixed-quota scheduling scope that has NO burst allowance together with a bursty multi-process fleet (spawning subagents, build workers, test runners): the fleet's demand spikes consume whole enforcement periods and the interactive

process freezes with them — in bursts, progressively as the fleet grows (the §11.4.103/§11.4.183 parallel-stream mechanics make the fleet's demand spiky BY DESIGN, so the collision is structural, not incidental). The consuming project MUST (1) grant the shared scope a burst allowance sized for the fleet's measured spike profile, OR (2) isolate the interactive process in its OWN scope (or leave it unquoted where host policy allows), OR both. Quota values MUST be HOST-ADAPTIVE — computed from real detected capacity with a documented floor and cap (§12.11's dynamic-detection discipline; §11.4.6 never hardcode), never a fixed literal that silently strangles a many-core host (the forensic clause-(b) instance); and a quota change is verified by READING THE LIVE scope's values BACK (§11.4.108 runtime signature — the configured value is not the effective value until read back). Composes §12.6 (memory ceiling) + §12.12 (thread/RLIMIT_NPROC headroom) as the THIRD orthogonal host-exhaustion axis: abundant RAM and abundant thread headroom do NOT imply an un-throttled scheduler — each axis has its own authoritative instrument, and none substitutes for another.

Honest boundary (§11.4.6): throttle telemetry proves WHERE enforcement periods were lost — it does NOT prove throttling is the ONLY degrader (network RTT floors, in-process GC/event-loop load, and sink-side effects remain separately measurable; the founding investigation ranked them independently and REFUTED five candidate mechanisms with the same instrumentation rigor); and clause (B) bounds PLACEMENT + SIZING — it does not exempt the fleet from the §12.6/§12.8/§12.12 host-safety ceilings, and it never licenses removing a quota that host policy requires (widen/isolate/burst within policy, per §12).

Classification: universal (§11.4.17) — scheduler-quota throttling is a platform-general mechanism (cgroup v1/v2, container runtimes, VM hypervisors); the consuming project supplies its concrete scope layout, quota/burst values, sampling windows, and instruments per §11.4.35. Composes §11.4.6 / §11.4.24 (build-resource stats — the sampler seam this telemetry extends) / §11.4.102 (systematic-debugging — the throttle sample is Phase-1 fact-gathering, and the forensic investigation reached it only by refusing to stop at healthy averages) / §11.4.108 (read the live scope back) / §11.4.128(1) (observer-effect budget — the sampler itself stays lightweight) / §11.4.201 (authoritative-source measurement; the averages-only verdict is a false-null) / §12.6 / §12.8 / §12.11 / §12.12. Propagation gate CM-COVENANT-114-225-PROPAGATION (literal 11.4.225) + recommended gates CM-QUOTA-THROTTLE-TELEMETRY (every progressive/bursty-degradation investigation's verdict cites throttle-accounting DELTAS sampled over a measured

window under load + a quiet-phase control; an averages-only verdict on a progressive-degradation issue → FAIL) and CM-INTERACTIVE-SCOPE-NOT-QUOTA-STARVED (no interactive-latency-sensitive process is co-resident in a fixed-quota no-burst scope with a bursty fleet; quota defaults are host-adaptive with a live read-back, never fixed literals) + paired §1.1 mutations (record a progressive-degradation verdict citing only averages → the telemetry gate FAILs; pin a fixed low quota literal as the default on a detected many-core host, or co-locate the terminal server with the fleet in a no-burst quota scope → the isolation gate FAILs; strip the literal → the propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md §11.4.225.

Non-compliance is a release blocker regardless of context. No escape hatch — no

--averages-prove-healthy , --skip-throttle-telemetry , --fixed-quota-default-OK , --interactive-in-fleet-scope-OK , --no-burst-needed , --skip-quiet-phase-control flag exists.

§ . . — Evidence-class-at-closure + standing detection pressure: machinery presence does NOT predict whether a fix holds — the EVIDENCE CLASS at closure (runtime vs source) under real DETECTION PRESSURE does; prose does not bind, seams do (research-derived, — —)

Forensic anchor (genericised — the 2026-07-22 reopen/first-touch root-cause forensics; operator mandate 2026-07-22, verbatim: "do deep research on how we have for very long time so much reopened workable items!? Why the quality of delivered releases was so poor mostly? We mostly detect the same major problems on first touch when new releases are available, and this happens over and over again!"). The single strongest empirical result across four independent evidence lines (internal forensics; a 76-source external literature pass; one cross-project corpus WITHOUT tracking machinery; one corpus WITH exemplary machinery): **every fix whose done-claim was preceded by an on-target RED → GREEN polarity flip that held recorded ZERO reopens; every fix confirmed only at source-green bounced.** Cross-corpus: the runtime-evidence-closure corpus recorded median attempts-to-stick = 1 and a ~0.3% corrected reopen rate (its own analysis marks that a FLOOR under weak detection pressure),

while the source-green corpus recorded 52% — itself a proven undercount; the healthy corpus's ONLY bouncing fix family was closed by a wrong-layer echo-assertion (a unit test asserting a Noop echo — the discriminator's predicted loser, sign right); 69 commits in a two-month window moved claims at "PENDING-BUILD"; the last audited release entered manual QA with **1 of 25** in-scope items validated at the user-visible layer — the operator was the FIRST user-layer oracle most claims had ever met, so first-touch rediscovery of the same majors was the deterministic output of the pipeline shape, not bad luck; a pixel-layer defect was "verified" by a file-existence check plus five greps; and the corpus WITHOUT machinery shipped a "fully working" release hours before a total field brick with its whole hermetic (wrong-layer) suite green. The law this evidence selects: **machinery presence (trackers, guards, gates, anchors) is signal and custody, NOT prevention; what predicts whether a fix holds is the EVIDENCE CLASS AT CLOSURE (runtime vs source) under real DETECTION PRESSURE. Prose does not bind — seams do. Coverage has a blind complement at any maturity.**

The mandate (ALL hold; the STRICT FORMALIZATION of §11.4.115(F)'s evidence-class clause + §11.4.146(D3)'s custody chain + §11.4.108's four layers — it supplies the closed taxonomy, rank floor, machine fields, and freshness contract those anchors did not state):

(1) Closed evidence-class ranking with a per-defect-layer floor, enforced at every closure seam. Defect layers form the closed set {user-visible, runtime, artifact, source} (§11.4.108); evidence classes form the closed RANKED set {runtime > artifact > source}. Every terminal or ready-class status write carries evidence whose class meets its defect layer's floor: a user-visible or runtime-layer defect can NEVER close on artifact- or source-class evidence; source-on-source stays legal (the §11.4.201(1) false-refusal guard — legitimate source-layer work is never refused).

(2) Class is proven by MACHINE FIELDS, never by a label. `runtime` requires a target fingerprint read FROM the target at run time (§11.4.115(F)) plus a runtime observable; `artifact` requires the artifact path plus its content hash; `source` requires the source ref. A class label without its fields is SHAPE-INCOMPLETE and refused; unreadable evidence is BLIND, never accepted.

(3) The ANTI-ECHO rule. A "runtime observable" that IS a grep/source-scan transcript is WRONG-LAYER and refused — a grep can never observe a pixel; the wrong-layer oracle dies by construction (the pixel-defect-by-five-greps closure and the Noop-echo family both cross this seam and are refused at it).

(4) Per-status evidence preconditions. Every status in the §11.4.15/§11.4.33 lifecycle carries a CHECKED evidence precondition — not only the terminal write. Forensic FACT: a penultimate "ready"-class status functioned as design-complete-but-UNVERIFIED and parked 79 items — 42% of the done-claiming population; the operator sampling "done" items was sampling a set dominated by never-verified work without either party being wrong about their own definition (semantic drift between claim-maker and claim-consumer). A status that sounds like progress but carries no checked precondition becomes the parking lot for unverified work.

(5) Fix-kind honesty. Every defect closure classifies its fix-kind; a STATE-ONLY repair (restart / cache-clear / re-deploy / manual re-run) or a mitigation CANNOT close a defect item as fixed — it is recorded as a mitigation and the item stays open or is honestly reclassified (composes §11.4.102's Iron Law + §11.4.146; forensic FACT: a sibling corpus's recurrence chains — 11 chains, median 3 and max 7 attempts per defect — were dominated by state-only "fixes" that closed and bounced).

(6) STANDING DETECTION PRESSURE. Detection pressure is a CAUSAL INPUT to every measured quality rate — a reopen/defect/quality figure reported without its detection-pressure regime is incomparable and misleading (§11.4.6): the measured 52%-vs-~0.3% cross-corpus split is explained by pressure + closure class, not machinery, and the healthy corpus's one genuine recurrence surfaced at 24 days — the next SWEEP, not the next use ("we only see what we test", §11.4.118). Therefore: every registered, topology-present guard carries a FRESHNESS CONTRACT — a verdict on the CURRENT artifact fingerprint, no older than a declared staleness budget; the never-executed / stale-fingerprint / over-budget set feeds a STANDING, risk-ordered re-run queue (most-reopened-first per §11.4.189, then stalest-first) that the zero-idle loop (§11.4.94/§11.4.97) drains as regular development-cycle work; the §11.4.135 release seam remains the blocking BACKSTOP, never the only executor. Forensic FACT: two standing guards, finally executed during the forensics, immediately emitted the FAILs that had been latent the whole time — one reproduced the operator's loudest defect 3/3; registration had been treated as coverage while execution was nobody's job. Topology-absent

guards are enumerated, never queued (no false pressure, §11.4.201(1)); an empty registry is BLIND, never "all fresh"; staleness budgets are consumer data (§11.4.35) bounded by the alarm-fatigue trade-off — the queue is a work FEED, only the release seam blocks.

Regular-cycle + review + pull binding (operator mandate 2026-07-23). This anchor applies during EVERY regular development cycle (the freshness queue is standing §11.4.94 work, and every closure crosses seam (1)); it is CHECKED through every code-review per §11.4.194(6); and on every constitution fetch/pull the §11.4.32 post-pull validation sweep + the §11.4.164 auto-propagation hook analyze the new additions and run their gates against the post-pull tree, with violations tracked (§11.4.15) and cleared — no parallel mechanism is invented.

Honest boundary (§11.4.6). (i) The discriminator is a measured correlation with a stated mechanism and zero counter-cases in the record — NOT a randomized trial (guarded items had less operator exposure; the clean exemplars are few; the confound is carried, not hidden). (ii) Fabricated machine fields are not fully closed here — provenance stays with §11.4.115(F) harness-written verdicts + the §11.4.135 candidate-fingerprint join, where an invented fingerprint fails to match the real candidate and reads as absence; layered, not eliminated. (iii) Oracle calibration stays §11.4.107(10) — runtime-class evidence from a mis-calibrated analyzer still passes this seam; goldens are necessary, not sufficient. (iv) The defect-layer taxonomy is authored data — mislabelling a user-visible defect as source-layer lowers its floor; the §11.4.194(6)(a) review class is the detective counter (cheap to spot: the item's own description names user symptoms). (v) Re-running a weak guard often proves nothing — oracle strength remains §1.1/§11.4.115(F)'s job, and orthogonal discovery (§11.4.85 chaos, §11.4.118 discovery pressure) is a different feed this queue does not generate.

Reference mechanisms (proven, NOT wired — §11.4.6/§11.4.205). Runnable seam mechanisms with RED-first POCs, golden-good/golden-bad/negative-control fixtures (§11.4.107(10)), live-data runs, and stated boundaries exist in this repository's quality-solutions research corpus (docs/research/quality/solutions/ — SOL-01 status-custody, SOL-02 verdict-coverage, SOL-04 evidence-class, SOL-09 detection-pressure; 65/65 GREEN across the ten-POC set on a deterministic second iteration). Per §11.4.205 they carry NO force until their seams are wired by separate tracked work items (§11.4.197); §11.4.227(A)'s ledger is the instrument that keeps that wiring debt visible.

Classification: universal (§11.4.17) — no hardware/vendor/project literal; the consuming project supplies its defect-layer bindings, evidence-store paths, verdict registries, and staleness budgets as DATA per §11.4.35. Composes §11.4.1 / §11.4.5 / §11.4.6 / §11.4.15 / §11.4.33 / §11.4.69 / §11.4.94 / §11.4.97 / §11.4.102 / §11.4.107(10) / §11.4.108 / §11.4.115(F) / §11.4.118 / §11.4.123 / §11.4.132 / §11.4.135 / §11.4.146(D3) / §11.4.185 / §11.4.189 / §11.4.194(6) / §11.4.201 / §11.4.205 / §11.4.214. Propagation gate `CM-COVENANT-114-226-PROPAGATION` (literal `11.4.226`) + recommended gates `CM-EVIDENCE-CLASS-AT-CLOSURE` (every terminal/ready-class status write carries class-matched, machine-field-proven evidence meeting its defect layer's floor; a grep-transcript "runtime observable" → FAIL; a state-only repair closing a defect as fixed → FAIL) and `CM-GUARD-FRESHNESS-SCHEDULER` (every registered, topology-present guard is fresh-on-the-current-fingerprint or present in the visible, risk-ordered re-run queue; an empty registry → BLIND, never "all fresh") + paired §1.1 mutations (close a runtime-layer item on source-class evidence → the closure gate FAILs; attach a runtime label with no machine fields → FAILs; let a registered guard sit past its staleness budget with no queue entry → the scheduler gate FAILs; strip the literal → the propagation gate FAILs; gate-code = separate work item, NOT claimed shipped §11.4.6).

Canonical authority: constitution submodule `Constitution.md` §11.4.226.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--source-green-closes-runtime-defect`, `--label-is-class`, `--echo-observable-OK`, `--registration-is-coverage`, `--rate-without-pressure-regime`, `--mitigation-closes-defect`, `--skip-freshness-queue` flag exists.

§ . . – Governance-corpus self-custody: the rule corpus is bound by the same custody it imposes – every named gate is implemented or registered-deferral under a monotone ratchet, and propagation gates count anchor BLOCKS (exactly-once + lockstep-identical), never bare literals (research-derived, – –)

Forensic anchor (measured on THIS corpus, 2026-07-22 — commands and control needles preserved in the root-cause analysis). The canonical rule file held 10,735 lines and 220 distinct §11.4.N anchors; **413 named CM-* gates, of which 241 (58%) had no implementation in either canonical gate site, plus 85 literal "gate-code = separate work item" deferrals;** the status-custody seam DESIGNED on 2026-07-17 was still prose on 2026-07-22 — the corpus's highest-value remediation fell into its own gap within five days; the two anchors that state the fix-confirmation discipline verbatim predate the failures they failed to prevent by five-plus weeks. The positive complement: **every fully-mechanised rule in the record either prevented or exposed its failure mode** (the suite PENDING-semantics + the release verdict-coverage seam landed within a day of diagnosis and now block) — the gap is not rule quality, it is the prose-to-seam conversion rate. The SHAPE defect: the propagation-gate family — the ONLY gate family that runs everywhere — asserts literal-presence ≥ 1 , mechanically rewarding restatement over implementation ("the corpus measures its own health in words present, and so it grows words"); measured consequences: two anchor blocks duplicated verbatim-divergently in ALL FOUR mirrors while every presence-gate stayed green; one anchor NUMBER heading two DIFFERENT mandates (the §11.4.140/§11.4.141 collision — independently re-detected by the block-integrity instrument's live run; resolution is an operator-owned re-mint, tracked); two canonical entry-point mandates contradicting each other (the mandated resumption file unreachable from the mandated README). External corroboration: a mandated-checklist adoption WITHOUT workflow embedding produced no measurable improvement across 101 hospitals (the Ontario null, cited in the external-research corpus) — prescription without a seam does not move outcomes at any scale.

The mandate (ALL hold):

(A) NAMED-GATE LEDGER + IMPLEMENTATION RATCHET. Naming a gate is a commitment with a same-commit cost: every CM- * -class gate token named anywhere in the governance corpus MUST, at every build (of the governance repo and of every wired consumer), be either IMPLEMENTED — the token present in an EXECUTABLE gate site; prose carriers NEVER count (§11.4.201(7)(a)) — or covered by a REGISTERED DEFERRAL row pointing at a tracked work item (§11.4.197). Silent gate debt is refused at the build seam. The unimplemented count is a MONOTONE-DECREASING ratchet, validated per §11.4.201(8) against the definition of done (at the correct end-state — every named gate implemented — the count is 0, so the metric reaches target and is valid as a necessary floor). A gate NAME that vanishes from the corpus without an explicit removal citation FAILs — the metric-gaming channel (deleting names to lower the count) is closed structurally; repeals stay legal and visible (the §11.4.166 pattern). **An anchor's "done" state is its SEAM landing, not its TEXT landing** — a rule authored without its seam is, by §11.4.205's own criterion, in the worse-than-nothing state, and this ledger makes that state visible, counted, and shrinking instead of silent. Brownfield adoption (ratchet down from the measured baseline vs hard-fail day one) is an operator §11.4.66 decision mirroring §11.4.224(E)'s adoption fence — never an invented ratchet.

(B) ANCHOR-BLOCK INTEGRITY — the propagation-gate SHAPE fix. Every propagation-class gate MUST count anchor BLOCK-STARTS (structural, line-anchored heading patterns), never bare literals: (i) exactly ONE block per anchor per governance file — 0 is the old absence check, > 1 is the DUPLICATION check, FAIL naming the anchor and file; (ii) content-hash EQUALITY across the declared lockstep mirror set (§11.4.157) — divergent copies FAIL naming both files (compact-vs-full variance ACROSS layers — canonical file vs mirror — remains legitimate; equality binds only within the declared lockstep set, and canonical-vs-mirror semantic drift stays the §11.4.186 cross-doc family's job); (iii) mid-body citations of an anchor literal are CARRIERS and never count (§11.4.201(7)(a)); (iv) ZERO blocks extracted from a governance file is BLIND-OR-EMPTY, never clean (§11.4.201(6)); (v) one anchor NUMBER heading two different mandates is a COLLISION and FAILs (§11.4.54's never-reuse discipline applied to anchor ids); (vi) the block extractor MUST capture full dotted ids so a sub-anchor (...10.A) never prefix-matches its parent (...10) — a measured §11.4.201(1) false positive in the founding instrument itself, fixed RED-first.

(C) Consolidation pressure + honest boundary. Nothing mechanical decides SEMANTIC contradiction between two different anchors — only the decidable projections are mechanized here (block counts, lockstep hashes, anchor numbers; fenced verdict scopes stay §11.4.112(5), seam declarations stay §11.4.120); the remainder is §11.4.142/§11.4.194 review territory — stated, never claimed. This anchor adds the measured enforcement-ratio INSTRUMENT, not a growth ban: the counterweights to corpus growth remain §11.4.17 classification, §11.4.141 token-efficiency, and the visible-repeal precedent (§11.4.166).

Regular-cycle + review + pull binding (operator mandate 2026-07-23). The ledger regenerates and diffs on every governance edit and every build — regular development cycles, never an on-demand audit; a change that NAMES a gate is review-checked for its implementation-or-deferral per §11.4.194(6)(c); and on every constitution fetch/pull the §11.4.32 post-pull sweep + the §11.4.164 auto-propagation hook run the ledger + block-integrity checks against the post-pull tree, with violations tracked (§11.4.15) and cleared — no parallel mechanism is invented.

Reference mechanisms (proven, NOT wired — §11.4.6/§11.4.205). SOL-05 (gate ledger + ratchet) and SOL-06 (anchor-block integrity) under docs/research/quality/solutions/ are runnable RED-first POCs with golden fixtures: SOL-05's live read-only run reproduced the 413-named-gate census exactly (234 unimplemented by its slightly-laxer structure rule — the 7-gate methodological delta stated in-doc); SOL-06's live run confirmed the four mirrors CLEAN post-F7 and independently re-detected the §11.4.140/§11.4.141 collision. Per §11.4.205 they carry NO force until wired by separate tracked work items (§11.4.197).

Classification: universal (§11.4.17) — no project literal; every consuming project (and this governance repo itself) supplies its gate-site list, deferral registry, and lockstep mirror set as DATA per §11.4.35. Composes §11.4.6 / §11.4.15 / §11.4.17 / §11.4.32 / §11.4.54 / §11.4.66 / §11.4.75 / §11.4.141 / §11.4.157 / §11.4.164 / §11.4.166 / §11.4.186 / §11.4.194(6) / §11.4.197 / §11.4.201 / §11.4.205 / §11.4.224(E). Propagation gate CM-COVENANT-114-227-PROPAGATION (literal 11.4.227 — itself REQUIRED to be block-start-shaped per clause (B)) + recommended gates CM-GATE-LEDGER-RATCHET (every named gate implemented-or-registered-deferral; the unimplemented count monotone-decreasing; vanished names cite removals) and CM-ANCHOR-BLOCK-INTEGRITY (exactly-once per anchor per file + lockstep content-hash equality + carrier exclusion + BLIND-on-zero + collision detection) + paired §1.1 mutations (name a gate with no

implementation and no deferral row → the ledger gate FAILs; delete a gate name without a removal citation → FAILs; duplicate an anchor block in one mirror, let two lockstep copies diverge, or re-mint an existing anchor number → the integrity gate FAILs; strip the literal → the propagation gate FAILs; gate-code = separate work item, NOT claimed shipped §11.4.6).

Canonical authority: constitution submodule `Constitution.md` §11.4.227.

Non-compliance is a release blocker regardless of context. No escape hatch — no

```
--name-gate-without-seam , --silent-gate-debt-OK , --delete-name-to-  
11 4 228  
lower-count , --presence-gte-1-suffices , --duplicate-anchor-OK , --  
mirror-drift-OK , --reuse-anchor-number flag exists.
```

§ . . — Cross-agent extension
lifecycle mandate: per-platform compatibility
declaration, source tracking, per-extension
documentation with compatibility matrix, and
auto-wiring of constitution extensions into all
detected agent platforms on constitution pull
(research-derived from Extensions Catalog
forensics, — —)

Forensic anchor (FACT, 2026-07-24 — measured from the live Extensions Catalog at `docs/extensions/EXTENSIONS_CATALOG.md`). The project's extension ecosystem — skills, MCP servers, ACPs, LSPs, plugins, and supporting infrastructure (34 skills, 1 MCP, 38 provider aliases, 21 CCR providers, 4 daemons, plus a separate HelixSkills repo) — exhibited four structural gaps: **(G1)** of 18 active Claude Code project skills, only Claude Code and OpenCode platforms are explicitly supported — HelixCode, Gemini CLI, and Qwen Code have zero recorded loading paths despite being governed by this constitution's own lockstep carriers (§11.4.157 — `CLAUDE.md`, `AGENTS.md`, `QWEN.md`, `GEMINI.md`); **(G2)** 4 of 18 active skills are symlinks into the constitution submodule (correct pattern) while 14 are local-only directories with NO traceability record to which original source/master each belongs to — a git-archive snapshot of a skill without its origin is a fork with no master; **(G3)** the Extensions Catalog is a single snapshot document with NO per-extension per-platform compatibility matrix and NO mechanism to re-detect drift when extensions change; **(G4)** the §11.4.164 `post_update_hook.sh` detects changed governance MDs + `scripts/hooks/skills/mcp/plugins` and registers them,

but it does NOT auto-wire constitution skills into all detected agent platforms — 4 constitution skills (`media-validator` , `multitrack` , `scheduled-work-queue` , `session-sync`) remain un-wired on both Claude Code (not symlinked into `.claude/skills/`) and OpenCode (no loading path to `.opencode/skills/`), silently absent while their governance files claim they exist — the canonical "docs say present, runtime says absent" gap §11.4.108 forbids.

The mandate (ALL hold — four sub-clauses, one lifecycle):

(A) Cross-agent extension compatibility. Every extension governed by this constitution — Skill, MCP server, ACP bridge, LSP, plugin, provider alias, embedding, daemon, or any future extension class — MUST:

1. **Ship a per-platform compatibility declaration** listing every governed CLI agent platform at minimum: Claude Code, OpenCode, Gemini CLI, Qwen Code. The declaration states, per platform, one of the closed vocabulary: `SUPPORTED` (verified working with a recorded per-platform liveness probe), `UNTESTED` (present but not yet verified — honest gap, tracked per §11.4.197), `UNSUPPORTED` (platform genuinely cannot load this extension class — per §11.4.112 structural-impossibility, with cited evidence), or `CONFIG-ONLY` (declared in config, never verified as wired per §11.4.139).
2. **NEVER cause loops, hangs, crashes, or broken behavior** in any platform that attempts to load it. An extension that causes a platform malfunction is a §11.4.228(A) violation of PASS-bluff severity — it claims compatibility it does not provide, producing the exact "green suite but broken-for-user" failure mode §11.4 forbids.
3. **Cross-platform liveness testing is mandatory** — the §11.4.169 closed test-type set applies: the extension MUST ship a per-platform liveness test (does the agent load it + produce correct output?) that runs on every declared-`SUPPORTED` platform at minimum. A platform where testing is genuinely infeasible (no runtime available on this host) is an honest §11.4.3 SKIP-with-reason recorded in the compatibility declaration, NEVER a faked PASS.

(B) Extension source tracking. Every extension MUST be traceable to its exact original source:

1. **Canonical source record** — each extension carries machine-readable provenance: originating repository URL, commit SHA (or tag+delta description when `UNTRACED`), canonical file path within that origin. This record lives in an

`EXTENSION_SOURCE.yaml` in the extension's directory or in a consumer-maintained extensions registry. An extension with no provenance record is `UNTRACED` and a tracked §11.4.197 item — never silently trusted as correct.

2. **Symlinks preferred over copies** — a copy is a fork with no master; a symlink (or a git submodule pointer, or a reflink share) preserves the canonical source relationship and makes divergence mechanically detectable. Where a copy is genuinely required (platform cannot traverse symlinks, an out-of-repo directory), the `SOURCE` record **MUST** state the canonical origin + why a copy was necessary (§11.4.6 honest rationale). Blind copies without provenance are §11.4.124-class dead-code-in-waiting.
3. **Per-project extensions catalog required** — every governed project **MUST** maintain an extensions catalog (`docs/extensions/EXTENSIONS_CATALOG.md` or equivalent declared per §11.4.35) enumerating every active extension, its source, its per-platform compatibility status, and its last-verified timestamp. The catalog is §11.4.86 roster/corpus-backed (sha256 fingerprint of the sorted extension keyset, NOT mtime), re-syncs out-of-the-box on any extension change. The catalog is the **SINGLE SOURCE OF TRUTH** for "what extensions exist" — ad-hoc per-agent skill directories without catalog entries are `UNCATALOGUED` and a tracked §11.4.197 item.

(C) Mandatory extension documentation with per-platform compatibility matrix. The extensions catalog **MUST**:

1. Carry a **per-platform compatibility matrix** (rows = extensions, columns = governed platforms, cells = compatibility status from the closed vocabulary in clause (A)(1)) — readable at a glance. A cell with no recorded status is `UNTESTED` by default, never silently assumed compatible.
2. Be maintained per §11.4.65 (four-format export: `.md` , `.html` , `.pdf` , `.docx`), updated on every extension add/change/removal, with §11.4.44 revision header.
3. Include **liveliness testing evidence per platform** — per extension × platform, a per-platform probe run result (load-ok / load-failed-for-reason / skip-with-reason) with a timestamp and evidence path per §11.4.69 `feature_class=extension_loading` . A platform declared `SUPPORTED` with no recorded probe is `UNTESTED` overridden — a classification mismatch and a §11.4.6 violation.

4. Auto-sync via the §11.4.106 docs-chain engine or equivalent — a stale catalog where extension drift has occurred (catalog says SUPPORTED, extension was removed or renamed) is a §11.4 PASS-bluff at the extension-management layer.

(D) Auto-wiring of constitution extensions into all detected agent platforms on constitution pull (§11.4.228, STRENGTHENS §11.4.164). The §11.4.164

`post_update_hook.sh` seam is EXTENDED — it MUST, after detecting changed extension artifacts in the constitution submodule:

1. **Detect all governed agent platforms** on the consuming host (Claude Code via `.claude/`, OpenCode via `.opencode/` or `opencode.json`, Gemini CLI via `~/.gemini/` or project `GEMINI.md`, Qwen Code via `.qwen/` or project `QWEN.md` — the platform-specific detection logic is consumer DATA per §11.4.35, never a hardcoded list). A platform whose detection mechanism is absent is an honest `DETECTION_UNSUPPORTED` gap recorded in the catalog, NOT a reason to skip the remaining platforms.
2. **Wire constitution skills into each detected platform** — for each detected platform, for each constitution skill not already present:
 - Claude Code: symlink from `.claude/skills/<skill-name>/` → `constitution/skills/<skill-name>/`
 - OpenCode: symlink from `.opencode/skills/<skill-name>/` → `constitution/skills/<skill-name>/`
 - Gemini CLI, Qwen Code: per-platform equivalent wiring mechanism (consumer-supplied DATA, §11.4.35)
3. **Register new/changed MCP servers** into each platform's MCP configuration (`.mcp.json`, `opencode.json`, equivalents) — per the existing §11.4.164 MCP-registration seam, applied to every detected platform.
4. **Self-validate** — after wiring, probe each platform that supports a `--list-skills / --list-extensions` equivalent or fall back to a filesystem check (skill/MCP directory present + parseable). A failed wiring is logged as a §11.4.32 post-pull finding — never silently absorbed.
5. **The hook is idempotent** — re-running it does not duplicate symlinks, corrupt configs, or lose existing non-constitution extensions. Every operation is dry-run-capable (report what WOULD change without changing it).

6. **Honest boundary (§11.4.6)** — auto-wiring makes constitution extensions REACHABLE to each platform; it does NOT guarantee each platform's runtime ACCEPTS them (a platform may have a different skill format, an incompatible MCP protocol, or a missing runtime — all honest UNTESTED / UNSUPPORTED gaps recorded in the catalog, never faked pass), and it does NOT make the extensions correct (§11.4.108 + the extension's own tests). The wiring gap this clause closes is the REACHABILITY gap — a skill that exists in the constitution but is invisible to a governed platform has the gap §11.4.108 forbids (source-present, runtime-absent).

Composes with §11.4.164 (post_update_hook — the seam this anchor extends) / §11.4.157 (GEMINI.md lockstep — the platforms whose extensions this governs) / §11.4.65 (universal export — the catalog's four formats) / §11.4.86 (roster/corpus-backed auto-sync — the catalog's fingerprint discipline) / §11.4.106 (docs-chain — the catalog's sync engine) / §11.4.69 (sink-side evidence — feature_class=extension_loading) / §11.4.124 (dead-code — blind copies are dead-code-in-waiting) / §11.4.3 (SKIP-with-reason — untestable platforms) / §11.4.112 (structural impossibility — UNSUPPORTED platforms) / §11.4.139 (clean artifact — CONFIG-ONLY un-wired class) / §11.4.32 (post-pull sweep — wiring failures) / §11.4.108 (runtime signature — a platform load is the extension's runtime signature) / §11.4.197 (research completion — per-platform wiring gaps are tracked items) / §11.4.169 (test types — per-platform liveness tests) / §11.4.142 / §11.4.194 (code-review — extension changes cross the same quality gauntlet). Classification: universal (§11.4.17) — no hardware/vendor/project literal; the consuming project supplies its platform detection mechanism, per-platform wiring commands, extension registry path, and catalog path as DATA per §11.4.35.

Propagation gate CM-COVENANT-114-228-PROPAGATION (literal 11.4.228).

Recommended gates:

- CM-EXTENSION-PER-PLATFORM-COMPAT (every governed extension carries a per-platform compatibility declaration using the closed vocabulary; a declared-SUPPORTED platform with no recorded liveness probe → FAIL)
- CM-EXTENSION-SOURCE-TRACKED (every extension carries provenance with origin repo + commit SHA or explicit UNTRACED rationale tracked per §11.4.197; a blind copy without provenance → FAIL)

- `CM-EXTENSION-CATALOG-DOCUMENTED` (the extensions catalog exists, carries a per-platform compatibility matrix, is four-format-exported per §11.4.65, and its sha256 fingerprint is fresh per §11.4.86; a stale catalog with extension drift → FAIL)
- `CM-CONSTITUTION-EXTENSIONS-AUTO-WIRED` (the §11.4.164 `post_update_hook.sh` auto-wires constitution skills into every detected governed platform, idempotently, and self-validates; a constitution skill not wired into a detected `SUPPORTED` platform → FAIL, with explicit `DETECTION_UNSUPPORTED` / `UNSUPPORTED SKIP-with-reason` — the §11.4.201(1) false-positive guard)

All paired §1.1 mutations (gate-code = separate work item, NOT claimed shipped §11.4.6).

Canonical authority: constitution submodule `Constitution.md` §11.4.228.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--skip-per-platform-compat`, `--single-platform-OK`, `--extension-without-provenance`, `--copies-over-symlinks-OK`, `--skip-extension-catalog`, `--catalog-format-optional`, `--skip-auto-wire`, `--manual-wiring-suffices`, `--untraced-extensions-OK` flag exists.

§ 11.4.229 **– Live in-session task/todo tracker MUST always be up to date and fully in sync with the real work state – never stale, never showing completed work as pending nor pending as done (User mandate, – –)**

Verbatim operator mandate (2026-07-25): "make sure it is ALWAYS up to date fully in sync! CRITICAL: This MUST BE one of mandatory rules / constraints! Add it with all relevant details into the constitution Submodule and relevant related files."

The agent's LIVE IN-SESSION TASK/TODO TRACKER — the harness-provided task list / `TodoWrite` -class todo tool / any equivalent moment-to-moment execution-state surface the runtime exposes DURING a session — MUST ALWAYS be kept up to date and FULLY IN SYNC with the real live work state: never stale, never showing completed work as pending, never showing pending / blocked / abandoned work as done or in-progress. It MUST be updated the MOMENT any

tracked task's state changes — the instant a task is started (→ in-progress), completed (→ done), blocked, abandoned, split, re-scoped, or newly discovered — not batched to a later checkpoint, not reconstructed from memory at the end of the work, not left reflecting a superseded plan. A live task tracker that lags the real work is itself a §11.4 PASS-bluff at the execution-transparency layer: it reports a work state the agent is not actually in, and an operator (or a resuming agent / conductor reading it per §11.4.116) then makes decisions on a false picture of what is done, what is in flight, and what is owed.

(A) EVERY STATE-CHANGE UPDATES IT, IMMEDIATELY. The moment a task starts, completes, changes status, or is discovered, the tracker is updated to match — exactly one task marked in-progress at a time where the surface models that single-active-task convention, no completed task left marked pending, no pending task left marked done, no ghost task lingering for work that was dropped (dropped work is closed with a reason, never silently left "in-progress"). Deferring the update ("I'll mark it all done at the end") is precisely the staleness this anchor forbids.

(B) NEVER STALE, NEVER INVERTED (§11.4.6). The tracker's content is CAPTURED FACT about the real work state — never a guess, never aspirational, never a plan the actual work has since diverged from. Showing completed work as still-pending, or showing pending / blocked / abandoned work as done or in-progress, is a §11.4 / §11.4.1 bluff of PASS-bluff severity (falsely-done) or FAIL-bluff severity (falsely-pending); a tracker that has drifted from reality is corrected the MOMENT the drift is observed, never carried forward.

(C) DISTINCT ARTIFACT — COMPLEMENTS, NEVER REPLACES. The live in-session task/todo tracker is the EPHEMERAL, moment-to-moment execution-state surface WITHIN a single session — a DIFFERENT artifact from each of: §12.10's version-controlled CONTINUATION document, §11.4.131's standing committed session-resumption FILE, §11.4.97's operator-facing milestone progress-update cadence, §11.4.106(F)'s doc/DB sync enforced at the commit / build / constitution-pull WRITE-SEAMS, and §11.4.202 / §11.4.208's workable-item DB and operator-request-history ledgers. Those anchors bind COMMITTED / CROSS-SESSION / OPERATOR-REPORT / TRACKED-WORK artefacts; NONE of them binds the live in-session task/todo tracker itself — §11.4.229 closes exactly that gap. It COMPLEMENTS them (a correct live tracker is the truth the §12.10 / §11.4.131 /

§11.4.202 / §11.4.208 artefacts are materialised FROM), never substitutes for any of them, and its always-in-sync discipline is the in-session analogue of §12.10's always-current mandate applied to the harness task list.

(D) HONEST BOUNDARY (§11.4.6). Where the runtime exposes NO such live task-tracking surface, this anchor is honestly N/A for that runtime (§11.4.3 SKIP-with-reason — the discipline BINDS the moment such a surface exists, per the §11.4.96 latent-binding pattern); it mandates that WHATEVER live tracker the runtime does expose is kept in-sync, not that a particular tool must exist. "In sync" means the tracker reflects the real work state as the agent knows it at that moment — the captured fact of what has started / completed / changed — not omniscience about work the agent has not yet observed.

Classification: universal (§11.4.17) — a platform-neutral execution-transparency discipline reusable by ANY agent / runtime that exposes a live in-session task/todo tracker; the consuming runtime supplies its concrete tracker surface (harness task list / `ToDoWrite` / equivalent) per §11.4.35. Composes §12.10 (always-current CONTINUATION — §11.4.229 is its live-in-session-tracker analogue) / §11.4.131 (standing session-resumption file) / §11.4.97 (progress-update cadence — the operator-report sibling) / §11.4.106(F) (write-seam doc/DB sync) / §11.4.126 (default autonomous-loop — the loop's live task state IS the tracker) / §11.4.116 (real-time conductor ↔ framework sync — a resuming conductor reads the live tracker) / §11.4.202 / §11.4.208 (workable-item / request-history materialisation) / §11.4.6 (no-guessing — the tracker is captured fact, never aspirational) / §11.4.1 (falsely-pending is a FAIL-bluff). Propagation gate `CM-COVENANT-114-229-PROPAGATION` (literal `11.4.229` across the consumer fleet) + recommended gate `CM-LIVE-TASK-TRACKER-IN-SYNC` (where the runtime exposes a live in-session task/todo tracker, its state matches the real work state at every state-change — no completed-task-marked-pending, no pending-task-marked-done, no ghost/stale entry, no superseded-plan drift; honest §11.4.3 SKIP-with-reason where the surface is genuinely absent — the §11.4.201(1) false-positive guard) + paired §1.1 meta-test mutation (mark a completed task as still-pending, OR a pending/abandoned task as done, OR leave the tracker reflecting a superseded plan → the gate FAILs; strip the literal → the propagation gate FAILs; gate-code = separate work item, NOT claimed shipped §11.4.6).

Canonical authority: constitution submodule `Constitution.md` §11.4.229.

Non-compliance is a release blocker regardless of context. No escape hatch — no

```
--tracker-may-lag , --batch-todo-updates , --stale-task-list-OK , --  
1 1 4 2 3 0  
reconstruct-todos-at-end , --skip-task-tracker-sync ,  
--pending-may-be-done , --completed-may-stay-pending flag exists.
```

§ . . — **Parallelized-pipeline methodology: build validate overlap + affected-stages-only re-run + parallel multi-device testing/recording + fan-out-subagents-always (research-derived, - -)**

Forensic motivation (genericised): the sequential pipeline shape — build FULLY completes → THEN validation begins → on ANY discovery, restart the WHOLE cycle from scratch → all devices tested one-at-a-time — serialises the operator's wall-clock onto build latency + one-device-at-a-time evidence gathering, so a multi-hour containerised build followed by serial per-device validation and a from-scratch restart-on-discovery makes time-to-a-manual-QA-capable-release a multiple of what the genuinely-parallel work requires. §11.4.58's PWU pipeline already overlaps Stage 1 of round N+1 with Stages 4–5 of round N; §11.4.89 already backgrounds long tests; §11.4.103 already mandates ≥3 auto-backfilled parallel streams; §11.4.119 already partitions shared hardware by single-owner; §11.4.128/ §11.4.144 already mandate always-on recording + availability-following. §11.4.230 BINDS those into one operating discipline and adds the three invariants they did not each state on their own: (A) the build ↔ validate CONCURRENCY seam + affected-stages-only re-run-on-discovery, (B) all-devices-concurrent + worker-pool recording/processing, (C) fan-out-subagents as the standing default. It is the parallelization capstone the way §11.4.103 is the continuous-parallel-stream capstone — a strengthening bind, never a weakening.

(A) PIPELINE-OVERLAP PARALLELIZATION. Validation/verification work MUST run CONCURRENTLY WITH the build wherever a validation stage does NOT depend on the build's own not-yet-produced artifact — never deferred until the build finishes when it could have run alongside it. Deploy the moment the build is done (build-completion is the deploy trigger, per §11.4.121 no-commit-while-build-writes-tracked-artifacts + §11.4.108 artifact-must-exist ordering — the deploy waits for the artifact, the deploy-independent validation does not). Concretely, WHILE a build runs, the following run in parallel per the §11.4.96 SAFE-during-build catalogue:

pre-build source-layer gates + meta-test mutation sweeps + coverage/clash checks (§11.4.110), the §11.4.125/§11.4.142/§11.4.194 independent code-review on the batch's SOURCE, doc/DB sync (§11.4.106), and on-device validation of the PRIOR build (device-independent of the in-flight build, §11.4.119). On any NON-huge-blocker discovery, the cycle MUST: (1) bump the version code (a new, distinctly-identifiable candidate artifact — never re-use the failing artifact's id, §11.4.111 stable-name + §11.4.200 target-identity); (2) fix at the proven root cause (§11.4.102 systematic-debugging first, §11.4.115 RED-on-the-broken-artifact); (3) RE-RUN ONLY THE AFFECTED STAGES, PARALLELIZED — the stages whose inputs the fix changed (its own targeted validation FIRST per §11.4.130, then the affected regression set), NOT a serial restart-from-scratch of the entire cycle. The explicit efficiency target is a **4–5× reduction in time-to-a-manual-QA-capable-release** versus the fully-sequential shape — a measured target, tracked as build-and-validate wall-clock deltas (§11.4.24 build-resource stats), never an assumed one (§11.4.6). **HONEST BOUNDARY (§11.4.6 / §11.4.108 / §11.4.120):** the CURRENT build's ARTIFACT-layer (§11.4.108 layer 2) and RUNTIME-ON-CLEAN-TARGET-layer (layer 3) validation CANNOT precede its own artifact — a gate whose precondition is the build's output belongs at a post-build seam, never overlapped onto authorship (§11.4.120 seam-placement; §11.4.69

`artifact_not_yet_built`). "Concurrent with the build" therefore means the deploy-INDEPENDENT stages overlap; the CURRENT-artifact runtime validation begins the moment that artifact exists, in parallel across devices per (B). **"Affected-stages-only re-run" is bounded to NON-huge-blocker discovery:** a HUGE BLOCKER (release-blocking-severity defect / regression invalidating the cycle / blast radius reaching the batch's other fixes) still triggers the §11.4.129 STOP → fix-all → full-RESTART protocol (restart, never resume) — §11.4.230(A) does NOT weaken §11.4.129, §11.4.40 full-suite-retest-before-tag, §11.4.130 validate-the-fix-first, §11.4.132 risk-ordered validation, or §11.4.185 manual-QA-final; it narrows the from-scratch restart to exactly the cases those anchors require it, and parallelises everything else.

(B) PARALLEL MULTI-DEVICE TESTING + RECORDING/PROCESSING. Every available, reachable test device MUST be exercised CONCURRENTLY — one validation stream per device, never one-device-at-a-time when the devices are independent — each stream the SINGLE EXCLUSIVE OWNER of that device's exclusive resources per §11.4.119 (the owner drives; any concurrent stream targeting the same device is read-only), so per-device evidence is never cross-

contaminated (§11.4.119 evidence-integrity). Every device is ALWAYS-RECORDED per §11.4.128 (non-invasive, side-effect-free, deterministic layout) and AVAILABILITY-FOLLOWED per §11.4.144 (a dropped device logs an honest offline event + reconnect-and-resume, never a silent corpus hole, never a faked continuity). Recording CAPTURE + ANALYSIS/PROCESSING MUST be fanned out through a BOUNDED WORKER-POOL (a Go worker-pool or platform equivalent) so N devices' recordings are captured + analysed in parallel and per-device evidence never serialises onto a single processor — the pool bounded by §12.6 memory + §12.12 thread headroom. Per-device PASS cites its own captured-evidence artefact (§11.4.5/§11.4.69/§11.4.107) on a CLEAN/fresh deployment (§11.4.108 runtime-signature, §11.4.200 verify-deployed-artifact-identity-per-device); the highest-risk / most-reopened cases (§11.4.132/§11.4.189) get the deepest scrutiny FIRST on every device. **HONEST BOUNDARY (§11.4.6):** genuinely-shared exclusive resources (one HDMI/audio sink, one media-playback path) serialise per §11.4.119 single-owner-over-time even across devices; "all devices concurrent" means independent devices, never two owners of one exclusive resource.

(C) FAN-OUT-SUBAGENTS-ALWAYS. Subagent fan-out (§11.4.20/§11.4.70 subagent-driven; §11.4.187 multi-track ruler; §11.4.192 auto-backfill) is the STANDING DEFAULT for EVERY genuinely-parallelizable effort — batch investigation, test-authoring, gate-authoring, per-device validation (B), per-PWU development (§11.4.58), doc/research work — not a per-request opt-in and not reserved for "big" tasks. A parallelizable effort executed monolithically in the foreground when a fan-out path exists is a §11.4.230(C) violation (the §11.4.94/§11.4.97/§11.4.103 zero-idle discipline applied to HOW work is dispatched, not only WHETHER it runs). **HONEST BOUNDARY (§11.4.6 / §11.4.183):** "always fan out" is BOUNDED by (1) genuine parallelizability (contending/redundant/serially-dependent subagents are NOT fan-out — they are waste, and spawning them is NOT maximization per §11.4.183 maximum-USEFUL-throughput); (2) §12.6 60% memory ceiling; (3) §12.12 `RLIMIT_NPROC` thread headroom; (4) §11.4.58 ≤6 concurrent-agent cap; (5) §11.4.119 single-resource-owner (two subagents may not both drive one exclusive device); (6) critical-state sequencing (commits/pushes/tags stay serial per §11.4.88/§11.4.113). The bar is maximum USEFUL parallel throughput, never agent count; fan-out that exceeds a host-safety ceiling or contends a single resource is a violation, not compliance.

Classification: universal (§11.4.17) — a platform-neutral parallelization discipline reusable by ANY project that builds an artifact, validates it across multiple targets, and records device evidence; the consuming project supplies its concrete device set + serials, build/flash tooling, recording layout, and worker-pool implementation per §11.4.35. Composes §11.4.20 / §11.4.24 (wall-clock stats prove the 4–5× target) / §11.4.40 / §11.4.58 (the PWU pipeline this binds + extends) / §11.4.69 / §11.4.70 / §11.4.88 / §11.4.89 / §11.4.94 / §11.4.96 (SAFE-during-build catalogue) / §11.4.97 / §11.4.102 / §11.4.103 / §11.4.108 (artifact/runtime layer ordering — the honest boundary) / §11.4.110 / §11.4.115 / §11.4.119 (single-resource-owner — B's load-bearing partition) / §11.4.120 (seam-placement — A's honest boundary) / §11.4.121 / §11.4.125 / §11.4.126 / §11.4.128 (always-on recording — B's evidence source) / §11.4.129 (huge-blocker full-restart — the un-weakened carve-out) / §11.4.130 / §11.4.132 / §11.4.144 (availability-following — B's drop-handling) / §11.4.183 (maximum-USEFUL-parallelism — C's honest boundary) / §11.4.185 (manual-QA-final — the un-weakened terminal gate) / §11.4.187 / §11.4.192 / §11.4.200 (per-device deployed-artifact verification) / §12.6 / §12.12. Propagation gate `CM-COVENANT-114-230-PROPAGATION` (literal `11.4.230` across the consumer fleet) + recommended gates `CM-PIPELINE-OVERLAP-VALIDATION` (a build that runs with zero deploy-independent validation stages running concurrently when SAFE-catalogue work is queued → FAIL; a NON-huge-blocker discovery that restarts the whole cycle from scratch instead of re-running only the affected stages → FAIL) + `CM-PARALLEL-DEVICE-WORKER-POOL` (available independent devices validated serially instead of concurrently, OR recording capture/analysis serialised onto one processor when a worker-pool path exists → FAIL) + `CM-FAN-OUT-SUBAGENTS-DEFAULT` (a genuinely-parallelizable effort run monolithically in the foreground when a fan-out path exists AND all host-safety ceilings have headroom → FAIL) + paired §1.1 meta-test mutations (per §1.1; gate-code = separate work item, NOT claimed shipped §11.4.6).

Canonical authority: constitution submodule `Constitution.md` §11.4.230.

Non-compliance is a release blocker regardless of context. No escape hatch — no

```
--sequential-build-then-validate ,  
--restart-whole-cycle-on-discovery , --serial-device-testing , --  
serialise-recording-processing , --no-fan-out , --monolithic-  
parallelizable-work flag exists.
```


§ . . – Nano-precision model-tier selection: lightest-capable-model-first with dynamic mid-task escalation (User mandate, – – , CRITICAL)

Verbatim operator mandate: every work effort MUST use the LIGHTEST CAPABLE model, escalating tier only as complexity rises — a nano-precision, per-task model-tier choice, NOT "everything on the heaviest model." Forensic motivation (genericised, §11.4.6): dispatching every task — a trivial doc-sync, a one-line status read, a mechanical format fix — on the heaviest available model wastes capacity + token budget + wall-clock that a lighter tier would spend identically well, while the OPPOSITE failure (a hard multi-factor root-cause investigation dispatched on a tier too light to reason through it) produces exactly the shallow, guessed, or wrong-conclusion work §11.4.6 / §11.4.102 / §11.4.194 forbid. Neither "always heaviest" nor "always lightest" is compliant; the correct discipline is a PER-TASK, EVIDENCE-BASED tier choice that tracks the task's own complexity, re-evaluated as that complexity becomes known — the nano-precision selection this anchor mandates.

(A) LIGHTEST-CAPABLE-FIRST + ESCALATE-ON-COMPLEXITY. Every dispatched work effort — main-stream or subagent (§11.4.20/§11.4.70), single-track or multitrack (§11.4.187/§11.4.192) — is assigned the LIGHTEST model tier from the closed, capability-ordered ladder `haiku < sonnet < opus` genuinely sufficient to complete it correctly (the review/merge-conflict-resolution substrate is a SEPARATE, already-pinned tier per clause (E), not a fourth rung of this ladder). Tier selection is a DECIDABLE heuristic, not a guess (§11.4.6) — evaluated against, at minimum: (i) **task scope** (single mechanical edit / focused single-file change / multi-file-or-multi-component change); (ii) **reasoning depth** (a lookup-and-apply action vs. a multi-factor root-cause investigation per §11.4.102/§11.4.194); (iii) **blast radius** (a change touching one isolated file vs. one crossing shared state, hardware safety §11.4.133, or governance/gate seams §11.4.227); (iv) **novelty** (a previously-solved, pattern-matched task vs. a genuinely new problem with no established precedent). Closed-set tier bindings (illustrative, non-exhaustive — the consuming project supplies its own precise mapping per §11.4.35): `haiku` — trivial/mechanical work (simple greps/searches, doc-sync regeneration, status/field reads, format/lint fixes, boilerplate scaffolding); `sonnet` — moderate work (a focused single-component debug, a bounded single-file implementation, straightforward research or audit reads); `opus` — complex work (architecture/design decisions,

multi-factor reasoning per §11.4.194(1), cross-feature blast-radius analysis per §11.4.92 Pass 2/3, hard root-cause investigation per §11.4.102 that resists a first-pass hypothesis). A task whose complexity is genuinely UNDECIDABLE from available evidence starts at the NEXT tier UP from the best available estimate (the safe-reversible default per §11.4.101) rather than guessed downward.

(B) DYNAMIC MID-FLIGHT SWITCHING. Tier assignment is NOT a one-time, start-of-task decision — it is RE-ASSESSED continuously as the task's real complexity becomes known during execution. A task DISPATCHED as simple that PROVES harder mid-flight (an unexpected multi-factor interaction surfaces, a root cause resists the first two hypotheses, the blast radius turns out to cross a governance/hardware-safety seam) MUST ESCALATE to the next-heavier tier for its remainder — continuing on an under-powered tier past the point complexity is KNOWN is a clause-(D) false economy, not compliance. Conversely, a task DISPATCHED as heavy that DECOMPOSES into genuinely light, bounded, low-blast-radius subtasks (mechanical file edits following an already-designed pattern, doc regeneration, per-item status reads) dispatches THOSE subtasks on a LIGHTER tier via subagent fan-out (§11.4.20/§11.4.70/§11.4.183) rather than running every subtask on the parent's heavier tier by default — capability-matching applies at EVERY dispatch granularity, not only the top-level task.

(C) ALWAYS-ON COMPOSITION, NEVER A PER-REQUEST OPT-IN. Nano-precision tier selection is a STANDING DEFAULT mechanism, engaged automatically for EVERY work effort from the first prompt of a session onward (mirrors §11.4.126 default-autonomous-loop / §11.4.198 default-mechanisms-always-on) — applied TOGETHER WITH, never as a substitute for: §11.4.141 (heavy token-optimization — a lighter model dispatched with a bloated context is not the efficiency this anchor targets); §11.4.183 (maximum-USEFUL-parallelism — tier-matching each parallel stream, not merely running more streams on one fixed tier); §11.4.187/§11.4.192 (the multi-track ruler orchestration + continuous auto-backfill — each track's worker is tier-assigned per this anchor, not defaulted to one blanket tier project-wide); and §11.4.196/§11.4.198 (native-alias-first priority + always-on default mechanisms — this anchor governs WHICH TIER within whichever alias/account is selected, composing with, never replacing, §11.4.196's ALIAS-layer priority). None of these five composed anchors is weakened, duplicated, or re-specified by §11.4.231 — this anchor adds the TIER axis to their existing ALIAS/PARALLELISM/TOKEN axes.

(D) HONEST BOUNDARY (§11.4.6) — "LIGHTEST CAPABLE" MEANS

CAPABLE, NOT CHEAPEST-REGARDLESS. "Lightest capable" is a CONJUNCTION, not a discount: a tier choice that fails to actually complete the task correctly — producing a shallow investigation, a guessed root cause, an under-verified fix, or ANY form of bluff per the §11.4 anti-bluff covenant — is NOT "lightest capable," it is UNDER-POWERED, and under-powering a task that then fails, bluffs, or requires costly rework is a FALSE ECONOMY and a §11.4 violation of equal severity to any other PASS-bluff (the token/time saved at dispatch is spent many times over in re-investigation, re-review, and the risk of a shipped defect). Capability determinations are made from the clause-(A) decidable heuristic + observed task outcomes — EVIDENCE-BASED, never guessed (§11.4.6): "this always works fine on haiku " without a captured basis is exactly the guess §11.4.6 forbids, and a tier's genuine capability for a task class is established by the same anti-bluff discipline (§11.4.5/§11.4.69/§11.4.107) that governs every other claim in this constitution. A platform-chosen serving model outside the haiku < sonnet < opus ladder (e.g. fable served for general non-review work) is NOT a §11.4.231 dispatch-tier violation — the ladder binds the DISPATCHER's tier ASSIGNMENT, not the platform's own per-turn serving choice.

(E) HARD PINS OVERRIDE THIS DEFAULT FOR THEIR SCOPES. §11.4.231 sets the DEFAULT working tier for NON-review, non-merge-conflict-resolution work. It does NOT relax, override, or compete with the two existing hard model-substrate pins: §11.4.209 (every mandatory independent code-review — §11.4.125/ §11.4.142/§11.4.194 iterated to GO per §11.4.134 — runs on **Fable at xhigh effort, Opus at xhigh** on genuine Fable-unavailability, NEVER a lower tier or lower effort) and §11.4.211 (every main → feature/product/flavor merge-CONFLICT resolution runs on the SAME Fable-xhigh / Opus-xhigh substrate). Those two anchors are STRICTER, SCOPE-SPECIFIC pins that this anchor's general ladder does not reach — review and merge-conflict-resolution are NOT eligible for haiku/ sonnet/opus-by-complexity dispatch; they run on their pinned substrate ALWAYS, regardless of how "simple" the reviewed change or the conflict appears. A review or merge-conflict resolution dispatched on any tier other than the §11.4.209/§11.4.211 pin (absent a genuinely-proven Fable-unavailability fallback) is a violation of THOSE anchors, not a valid application of §11.4.231's general ladder.

(F) EFFORT-TIER SELECTION — the SAME nano-precision discipline applied to the reasoning-EFFORT axis (operator IMPORTANT mandate, 2026-07-26; verbatim: "Besides dynamic determining proper models to use, we MUST

decide the effort as well and switch between the effort settings same as we do now for chosen model! Labeling MUST BE extended for this details as well. Example: (T1/main - claude4 - opus - high)."). Every dispatched work effort is assigned not only a model tier per clauses (A)-(E) but ALSO the LIGHTEST reasoning-EFFORT setting from the closed ladder

`low < medium < high < xhigh < max` genuinely sufficient to complete the task correctly — chosen by the SAME clause-(A) decidable heuristic (scope / reasoning-depth / blast-radius / novelty), NOT a guess (§11.4.6). Illustrative bindings (non-exhaustive; the consuming project supplies its precise mapping per §11.4.35):

`low` / `medium` — trivial / mechanical work (doc-sync regeneration, status / field reads, format / lint fixes, boilerplate); `high` — moderate-to-complex reasoning (focused single-component debug, multi-factor analysis per §11.4.102 / §11.4.194); `xhigh` / `max` — the HARDEST verification / merge-conflict / architecture work.

(F.1) DYNAMIC MID-TASK ESCALATION — effort, like model tier per clause (B), is RE-ASSESSED continuously as the task's real complexity becomes known; a task dispatched at a lower effort that PROVES harder mid-flight MUST ESCALATE its effort for the remainder (continuing under-effort past the point complexity is KNOWN is the clause-(D) false economy), and a heavy task decomposing into light bounded subtasks dispatches THOSE at a lighter effort — effort-matching applies at EVERY dispatch granularity, not only the top-level task. (F.2) HONEST HARNESS BOUNDARY (§11.4.6, load-bearing — do NOT overstate the mechanism) — effort-tiering is REAL where the execution path exposes an effort control and a DOCUMENTED CAPABILITY-GAP where it does not: the current Claude Code **Agent tool** (subagent dispatch) exposes a `model` parameter but **NO effort parameter**, so a subagent's effort is NOT presently settable via that path (the §11.4.182 `<effort>` label field degrades honestly to `?` there); the

Workflow tool's `agent()` DOES take an `effort` argument, so effort IS settable there — the label reports whichever path is live, never fakes one. This is the SAME aspirational-vs-harness state the clause-(E) §11.4.209 / §11.4.211 Fable- `xhigh` pins already occupy: the pins mandate the effort, and whether a dispatch path can SET it is a harness capability reported honestly, never faked (§11.4.182). (F.3) HARD-PIN FLOOR PRESERVED — clause (E)'s hard pins are ALSO effort pins: every §11.4.209 mandatory independent code-review AND every §11.4.211 merge-conflict resolution runs at `xhigh` effort (Fable, or Opus at `xhigh` on genuine Fable-unavailability), and this general effort ladder NEVER tiers those two scopes below `xhigh` regardless of how "simple" the reviewed change or the conflict

appears — a review / merge-conflict resolution run below `xhigh` is a §11.4.209 / §11.4.211 violation, not a valid application of this ladder. (F.4) HONEST BOUNDARY (§11.4.6) — "lightest capable effort" is a CONJUNCTION, not a discount: an effort setting that yields a shallow / guessed / under-verified result is UNDER-POWERED, a FALSE ECONOMY, and a §11.4 violation of PASS-bluff severity, exactly as an under-powered MODEL tier is per clause (D); capability determinations are evidence-based, never guessed (§11.4.6).

Classification: universal (§11.4.17) — a platform-neutral model-capacity-matching discipline reusable by any project dispatching work across a tiered model roster; the consuming project supplies its own precise tier ↔ task-class mapping and its own model-name ladder per §11.4.35 (the `haiku / sonnet / opus` ladder and the Fable/Opus-xhigh pins are this project's current concrete instantiation, not a universally-fixed vocabulary — a consumer with a differently-named tier roster maps its own ladder onto the SAME lightest-capable-first + dynamic-escalation + always-on-composition + honest-boundary + hard-pin-override structure).

Composes §11.4.6 (no-guessing — capability determinations are evidence-based) / §11.4.20 / §11.4.70 (subagent-driven dispatch — the granularity at which tier-matching applies) / §11.4.92 (blast-radius analysis feeds the clause-(A) heuristic) / §11.4.101 (safe-reversible default on undecidable complexity — escalate up, never guess down) / §11.4.102 / §11.4.194 (reasoning-depth signals for the heuristic) / §11.4.126 / §11.4.141 / §11.4.183 / §11.4.187 / §11.4.192 / §11.4.196 / §11.4.198 (the always-on composition set, clause (C)) / §11.4.209 / §11.4.211 (the hard-pinned review/merge-conflict substrate this anchor does NOT touch, clause (E)).

Propagation gate `CM-COVENANT-114-231-PROPAGATION` (literal `11.4.231` present as a block-start, exactly-once per governance file, lockstep content-hash equality across the consumer fleet per §11.4.227(B)) + recommended gate `CM-MODEL-TIER-LIGHTEST-CAPABLE` (a dispatched effort whose logged/declared model tier exceeds the clause-(A) heuristic's indicated tier for a task later confirmed trivial-and-successful-at-the-lighter-tier is a WARN/candidate finding — never a hard FAIL on tier choice alone, since tier-adequacy is only provable in retrospect per clause (D); the gate instead HARD-FAILS the two decidable violations: (i) a review or merge-conflict-resolution effort NOT run on the §11.4.209/§11.4.211 pinned substrate, and (ii) a task whose complexity was PROVEN mid-flight [an escalation event was logged, or a downstream failure/bluff was traced back to it] that did NOT escalate tier before continuing) + paired §1.1 mutation (gate-code = separate work item, NOT claimed shipped — §11.4.6; mutation plan: (1) dispatch a

code-review labeled as running on `sonnet` instead of `Fable-xhigh` → gate MUST FAIL; (2) inject a logged mid-task complexity-escalation event with NO corresponding tier-escalation record → gate MUST FAIL; (3) golden-FALSE fixture: a task that stays on its initially-assigned tier throughout, with no escalation event logged → gate MUST NOT fire, per §11.4.201(1) the false-positive guard).

Canonical authority: constitution submodule `Constitution.md` §11.4.231.

Non-compliance is a release blocker regardless of context. No escape hatch — no

```
--always-heaviest-model , --always-lightest-model ,  
--skip-complexity-reassessment , --fixed-tier-per-session , --relax-  
review-pin , --relax-merge-conflict-pin , --per-request-tier-opt-in  
flag exists.
```

§ . . — Anti-mess long-op
orchestration: every long-op is registered,
single-owned-per-purpose, liveness-PROVEN,
handed-off-on-stop, safely-reaped, consistency-
checked, and escape-bounded — the single source
of truth for what is running, never inferred
(research-derived, — —)

Forensic anchor (genericised): in one cycle a verification sweep produced five distinct mess states — (1) multiple concurrent same-purpose sweeps ran unintentionally (no single-owner-per-purpose claim); (2) a sweep HUNG on a consumer-blocked-on-an-unclosed-producer pipeline and was undetected because liveness was inferred from process-alive (a §11.4.201(6) FALSE-NULL: a wedged op and a progressing op returned the identical "process alive" signal); (3) stopping the owning subagent SILENTLY KILLED its child long-op because ownership was implicit (no handoff-on-stop); (4) a stale/non-growing log was read as live progress (the same false-null, reader-side); (5) an un-completable tooling gate blocked a release for hours (no escape to a terminal state). Per §11.4.102 Phase 4.5 these are ONE architectural state-management flaw, not five bugs: **there was no single source of truth for what long-op is running, who owns it, its state, and its PROVEN liveness** — the §11.4.196(F) configured~~in~~-use gap generalised to project long-ops, so every mess state was un-observable, and an un-observable bad state is a §11.4/§11.4.1 bluff at the orchestration layer.

The mandate — a **long-op orchestration layer** binding every long-op (any op expected to exceed the runtime background-execution threshold per §11.4.89 — verification sweep, meta-test sweep, build, flash/deploy, test/QA suite, dispatched subagent), ALL of which hold:

(A) SINGLE SOURCE OF TRUTH — the long-op lifecycle registry. Every long-op MUST register BEFORE it runs and reach a terminal state (`complete` | `failed` | `reaped` | `handoff` | `blocked-escape`) when it ends, in a durable (never-tmpfs, write-temp-then-rename, §11.4.116 append-only event stream + atomic snapshot) registry recording `{op_id, purpose_key, owner, pid|pgid, cwd, state, last_heartbeat, heartbeat_seq, progress_offset, log_path, verdict, evidence_path}` . Success is read from the op's terminal verdict/ `result` , NEVER its process exit code. A long-op running with no registry record — or a record with no live op — is a FALSE-READY drift state.

(B) SINGLE-OWNER-PER-PURPOSE. No two same- `purpose_key` long-ops run at once; acquisition is a probe-race-safe claim keyed on `purpose_key` (§11.4.119). A live claim forces exactly ONE of REUSE (attach) or explicit SUPERSEDE (graceful handoff of the running owner, logged), never a silent second run.

(C) LIVENESS TRUTH-SOURCE. Progress is PROVEN, never inferred from process-alive: every long-op emits a monotonic heartbeat and/or an advancing progress offset; a watchdog reads the DELTA over a window and flags HUNG on no-advance past the op's declared no-progress budget (the §11.4.201(6)-(7)/ §11.4.89 discipline — `kill -0` is a necessary-not-sufficient pre-filter, a stale heartbeat/non-growing log is NEVER "progressing"). A consumer blocked on an unclosed producer write-end (§11.4.201(12)) is caught by the flat heartbeat regardless of internal state.

(D) OWNERSHIP-HANDOFF-ON-STOP. Stopping/killing an owner MUST hand its long-ops to the registry (`handoff` , re-adoptable per §11.4.147) or deliberately reap them (`reaped` , reason logged) — NEVER silently orphan or accidentally cascade-kill; long-ops launch detached (own process-group) so an owner stop does not cascade-signal them, and the stop path records each owned op's disposition. The loop done-condition is not satisfied while a `handoff` op is un-adopted.

(E) SAFE REAPING. An orphaned/hung op is reaped ONLY on PROVEN staleness — holder PID dead (`kill -0` fails) OR no-progress-budget exceeded per (C) — resolved by the process's real `/proc/<pid>/cmdline` (§11.4.196(D)/§12.12), never a bare `pgrep` /substring that could match a carrier or the reaper itself; a LIVE, ADVANCING op is NEVER reaped (§9.2); every REAPED/KEPT decision + resolved evidence is logged (§11.4.180).

(F) CONSISTENCY INVARIANT SWEEP. A fast invariant check runs before every gated transition (build/flash/tag) and on a standing cadence, asserting: no duplicate owner per purpose, no orphan, no stale lock (§11.4.180), no hung op (C), no registry-vs-reality drift (live-long-op set == registry live rows — the §11.4.196(F) check generalised). Any violation is surfaced (named, actionable) and the gated transition is REFUSED until resolved. The sweep is itself a guard (§11.4.201): golden-TRUE + golden-FALSE-with-carrier fixtures, conservative-safe REFUSE on an unresolvable signal, resolved evidence printed on every refusal.

(G) UN-COMPLETABLE-BLOCKER ESCAPE. A gate that cannot reach a terminal verdict MUST degrade, after a bounded budget, to an explicit evidence-backed decision — NEVER an indefinite silent block: §11.4.102 first (hung/(C) vs duplicate/(B) vs wrong-seam/§11.4.120 re-seat-at-activation); NOT-YET-RUNNABLE → §11.4.69 `artifact_not_yet_built` (non-blocking on an intermediate artifact, blocking only on the release candidate per §11.4.135); genuinely un-completable → §11.4.101 (autonomous reversible SKIP-with-reason + tracked §11.4.197 item, ELSE a bounded §11.4.66 operator decision) — the loop keeps progressing every other non-blocked item, it does not idle, and the escape NEVER fabricates a PASS.

Honest boundary (§11.4.6). This layer reduces orchestration mess; it does NOT make product bugs impossible, does NOT make a gate correct (§11.4.108/ §11.4.201 still bind each gate), and does NOT remove the operator from genuine blocks — it converts an indefinite silent stall into a bounded, evidence-backed, operator-visible decision. A heartbeat proves ADVANCE not CORRECTNESS; the registry proves what CLAIMS to run, (F) proves registry==reality. **Reuse-not-rewrite:** it is the multitrack per-track agent registry (§11.4.147/§11.4.116/§11.4.187) generalised from per-track workers to any long-op and bound through, per §11.4.196(F) — no parallel mechanism is invented.

Classification: universal (§11.4.17) — no hardware/vendor/project literal; consumers supply their long-op set, purpose keys, no-progress budgets, and stop-hook wiring as DATA per §11.4.35. Composes §11.4.89 / §11.4.101 / §11.4.116 /

§11.4.119 / §11.4.120 / §11.4.135 / §11.4.147 / §11.4.180 / §11.4.187 / §11.4.196(D) (F) / §11.4.201(1)(6)-(12) / §11.4.207 / §11.4.69 / §12.6 / §12.12 / §9.2. Propagation gate `CM-COVENANT-114-232-PROPAGATION` (literal `11.4.232`) + recommended mechanism gates `CM-LONGOP-REGISTRY-SINGLE-SOURCE` (A) / `CM-LONGOP-SINGLE-OWNER-PER-PURPOSE` (B) / `CM-LONGOP-LIVENESS-HEARTBEAT` (C) / `CM-LONGOP-OWNERSHIP-HANDOFF-ON-STOP` (D) / `CM-LONGOP-SAFE-REAP-PROVEN-STALE` (E) / `CM-ORCHESTRATION-CONSISTENCY-SWEEP` (F) / `CM-UNCOMPLETABLE-BLOCKER-ESCAPE` (G) + paired §1.1 mutations (run two same-purpose sweeps with no reuse/supersede → (B) FAILs; infer liveness from `kill -0` alone / drop the heartbeat-delta check → (C) FAILs; a stop that leaves an owned op with no terminal/handoff record → (D) FAILs; reap a live-advancing op OR reap by bare- `pgrep` substring → (E) FAILs; a live long-op with no registry row passes the consistency sweep → (F) FAILs; an un-completable gate blocks with no escape/decision → (G) FAILs; strip the literal → propagation FAILs; gate-code = separate work item, NOT claimed shipped §11.4.6). No escape hatch — no `--allow-duplicate-longop` , `--infer-liveness-from-pid` , `--silent-orphan-on-stop` , `--reap-live-op` , `--skip-consistency-sweep` , `--block-without-escape` , `--longop-without-registry` flag.

EXTENSION landed 2026-07-26 — §11.4.182 model-in-label + flawless alias/provider ↔ model consistency + real-time always-in-sync (User mandate, 2026-07-26). Verbatim operator mandates (2026-07-26): "Extend our logging to present models used since we will now have dynamic models choice and switching! From (example) (T1/main - claude4) into: (T1/main - claude4 - model_used), model MUST be dynamically obtained so we present name equivalents - short names - like: haiku, opus, sonnet, fable. However we MUST fully support all Claude Toolkit providers aliases, so real model used names MUST BE obtained in this case(s) as well! Examples: (T1/main - claude4 - haiku), (T1/main - claude4 - fable), (T1/main - claude4 - opus), (T1/main - claude4 - deepseek), (T1/main - claude4 - xiaomi) ... Labels MUST BE everywhere ALWAYS up to date in real time and in sync!" and "Adding model names into labels considers that providers names / aliases will be up to date too and in relation to model used since models belong to providers / aliases !!! Consistency MUST BE flawless!" This EXTENDS the §11.4.182 label discipline (it does NOT re-mint the anchor — the literal `11.4.182` and its propagation gate `CM-COVENANT-114-182-PROPAGATION` are UNCHANGED and stay GREEN; this carrier block cites §11.4.182, it is not a second §11.4.182 block-start per §11.4.227(B)). (1) **LABEL FORM EXTENDED** — the mandated §11.4.182

label is extended from (T<N>/<branch> - <alias>) to (T<N>/<branch> - <alias> - <model> - <effort>) (the trailing <effort> field added 2026-07-26 per §11.4.231 clause (F) — see the EFFORT-FIELD clause (7) below), <model> recording which model is CURRENTLY answering for that alias: for a native Claude alias, the SHORT name (opus / sonnet / haiku / fable , case-insensitive-substring-derived from the model id the agent platform reports for the most recent turn — DYNAMIC, obtained live at label-emission time, never a fixed/assumed value since native aliases have no persistent per-alias model and the serving model can legitimately switch mid-session per §11.4.231); for a claude-toolkit PROVIDER alias (deepseek , xiaomi , kimi-for-coding , opencode , ...) the REAL provider model id as declared in that provider's host-config registry (the SAME value the alias's own launch mechanism uses — a single source of truth, never re-derived). (2) **MODEL-FIELD DERIVATION (§11.4.6, never guessed)** — a closed two-case split on the SHAPE of the already-derived <alias> : (a) native-alias → read the model id from the CURRENT session's own transcript (the platform's per-project per-session transcript; LAST assistant-turn's reported model → short name); (b) provider-alias → read the model id from the host multi-account-toolkit provider registry (CMA_PROVIDER_MODEL). Either lookup that cannot resolve yields the honest ? — NEVER a fabricated or last-known value; an id obtained but matching none of the four native short-names is printed VERBATIM (real-but-unrecognised is more honest than a false ?). The 3-field legacy form (no <model>) remains a VALID, ACCEPTED label during migration — the enforcing guard-hook LABEL_RE admits both forms, and NEVER treats an honest ? model, or an absent model field, as a block-cause. (3) **FLAWLESS ALIAS/PROVIDER ↔ MODEL CONSISTENCY INVARIANT (operator NOTE, load-bearing)** — the <alias> / provider field and the <model> field MUST be MUTUALLY CONSISTENT AT ALL TIMES: the model shown MUST be a model that genuinely BELONGS to the alias/provider shown (models belong to providers/aliases), and the alias/provider name MUST itself be current — never a stale alias beside a fresh model, never a model that does not belong to the shown alias. This is guaranteed BY CONSTRUCTION, not by coincidence: the native model is read from THAT alias's own live session transcript, and the provider model is read from THAT provider's own registry entry — so the derived <model> is definitionally a model of the shown <alias> ; a derivation that ever emits a model not belonging to the shown alias is a §11.4.6 / §11.4.182 consistency violation, never a valid label. (4) **REAL-TIME, ALWAYS-IN-SYNC (§11.4.229/§11.4.106(F))** — the label is STATELESS and RE-COMPUTED

at EVERY emission (every human-invoked labeler run AND every synchronous PreToolUse guard-hook firing on every Agent/Task/TaskCreate dispatch) against the LIVE environment (current `CLAUDE_CONFIG_DIR` , current session id, current transcript tail, current provider `.env`) — so a mid-session model switch (`/model` , an autonomous §11.4.231 escalation) is reflected in the VERY NEXT emission with no cache to go stale: real-time BY CONSTRUCTION, not by a sync daemon.

Honest boundary (§11.4.6): the live label has NO persisted artifact that could drift, so the Docs-Chain (§11.4.106) sync engine — which propagates PERSISTED documents/DBs — does not apply to the stateless label itself; the operator's "use/extend Docs Chain for real-time responsiveness" ask attaches to any PERSISTED label-history / model-switch LOG surface (a §11.4.197 follow-up to be tracked), which, if adopted, IS §11.4.106-Docs-Chain-synced and §11.4.208-request-history-adjacent — never faked as already-shipped. (5) **HONEST SUBAGENT-INTROSPECTION LIMIT (§11.4.6/§11.4.197)** — the native-alias transcript lookup is proven reliable when invoked from the DISPATCHING process's own shell (the guard-hook's actual execution context at dispatch time); it is NOT proven reliable when a background-dispatched subagent introspects its OWN model from inside its own subprocess on hosts where the subagent shares its parent's session-transcript with no subagent-vs-conductor marker — recorded as an open limitation to be tracked, never silently claimed solved. (6) **DIVERGENT-LABELER RECONCILIATION (§11.4.227)** — the LIVE labeler wired via `.claude/settings.json` is `constitution/scripts/multitrack/track_branch_label.sh` (both it and the top-level `scripts/multitrack/` duplicate carry the model-field extension (functionally equivalent on the alias+model fields, not byte-identical — the constitution copy additionally carries the 2026-07-26 effort field the top-level copy does not yet mirror)); the two copies are reconciled to a single source of truth (top-level a thin wrapper exec-ing the constitution copy, §11.4.177 inherited-by-reference) as a reconciliation to be tracked, never a silent divergence. (7) **EFFORT FIELD (added 2026-07-26 per the operator IMPORTANT mandate — verbatim: "Besides dynamic determining proper models to use, we MUST decide the effort as well and switch between the effort settings same as we do now for chosen model! Labeling MUST BE extended for this details as well. Example: (T1/main - claude4 - opus - high).")**. The 5-field label appends `<effort>` after `<model>` , recording which reasoning-EFFORT setting is CURRENTLY in force for that work-stream, drawn from the closed set `{low | medium | high | xhigh |`

max} (the SAME effort ladder §11.4.231 clause (F) selects). (a) DERIVATION (§11.4.6, never guessed) — `<effort>` is read from whatever LIVE effort signal the execution path exposes (the Workflow-tool `agent()` `effort` argument, a future Agent-tool effort parameter, an env/config effort setting); a path exposing NO effort signal, OR any value outside the closed set, yields the honest `?` — EXACTLY like the `<model>` field's `?` fallback, NEVER a fabricated / assumed / last-known value. (b) HONEST HARNESS BOUNDARY (§11.4.6, load-bearing — do NOT overstate the mechanism) — effort-tiering is REAL where the execution path exposes it and a DOCUMENTED CAPABILITY-GAP where it does not: the current Claude Code **Agent tool** (subagent dispatch) exposes a `model` parameter but **NO effort parameter**, so a subagent's effort is NOT currently settable via that path and its `<effort>` field degrades honestly to `?`; the **Workflow tool's** `agent()` DOES take an `effort` argument, so effort IS settable + derivable there — the label reports whichever path is live, never invents one. This is the SAME aspirational-vs-harness state the §11.4.209 / §11.4.211 Fable- `xhigh` review/merge pins already occupy: the pins mandate the effort, and whether a given dispatch path can SET it is a harness capability the label reports honestly and NEVER fakes. (c) LEGACY-FORM ACCEPTANCE (migration) — the 3-field (`T<N>/<branch> - <alias>`) and 4-field (`T<N>/<branch> - <alias> - <model>`) label forms remain VALID and ACCEPTED; the guard `LABEL_RE` admits all three shapes and NEVER treats an absent `<effort>` field, or an honest `?` effort, as a block-cause (§11.4.6 — mirroring the §11.4.182 track/alias/model honest- `?` rule). (d) CONSISTENCY — the effort field is selected TOGETHER WITH the model by §11.4.231 (both tiers chosen for one task) and never contradicts it; a stale effort beside a fresh model, or an effort outside the closed set, is a §11.4.6 / §11.4.182 consistency violation, never a valid label. (e) REAL-TIME, ALWAYS-IN-SYNC (§11.4.229) — like `<model>`, `<effort>` is STATELESS and RE-COMPUTED at EVERY label emission against the LIVE environment, so a mid-task §11.4.231 effort escalation is reflected in the VERY NEXT emission with no cache to go stale. Propagation: `CM-TRACK-BRANCH-LABEL` is EXTENDED to also assert the labeler emits a parseable `<model>` field AND a parseable `<effort>` field, the guard `LABEL_RE` accepts the 3-field / 4-field / 5-field label forms (legacy 3-field and 4-field ALWAYS remain accepted during migration; a `?` effort is NEVER a block-cause), and the alias/provider ↔ model consistency holds; paired §1.1 mutation: strip the `<model>` derivation OR emit a model not belonging to the

shown alias → the extended gate FAILs (gate-code = separate work item, NOT claimed shipped §11.4.6). No new escape hatch — the existing §11.4.182 "no escape hatch" list applies to the extended label unchanged.

§ . . — Standing anti-mess control plane: a level-triggered, idempotent, invariant-catalogue reconciliation loop over the WHOLE persistent System state, orchestrating every per-domain enforcer, refusing gated transitions on un-reconciled drift (research-derived, - -)

Forensic anchor (genericised): every existing consistency enforcer in a multi-track / multi-agent / multi-submodule / multi-upstream System is EDGE-TRIGGERED — a PreToolUse hook fires on an action, a check runs at a seam (lock-acquisition / merge-time / build-time / commit-time), a gate runs when a build runs — so a mess that arises WITHOUT an edge firing is un-observed until it causes harm (two live examples captured 2026-07-26: a large uncommitted-batch pile-up sat unreconciled through a session, and multiple provably-stale VCS/wrapper locks were held by dead holders with no acquisition attempt to reap them — the §11.4.205 "a hook never installed / a check that runs only when asked" forensic generalised to the whole persistent state). Per §11.4.102 Phase 4.5 the ~19 persistent-state mess classes (Plane R repository: uncommitted pile-up, stale/duplicate locks, orphaned worktrees, half-applied cross-repo commits, drifted submodule pointers, partial merges, divergent mirrors, mis-placed work; Plane G governance: divergent mirrors, pointer drift, gate-ledger drift, anchor-number collision; Plane S tracking: doc/DB drift, cross-track SSoT divergence, ledger drift, handoff-doc drift; Plane P runtime: the §11.4.232 long-op set + un-reaped streams + dangling review state) are ONE architectural flaw — there is no LEVEL-TRIGGERED control plane that re-derives the actual persistent state and reconciles it against a declared clean state; an un-observed drift is a §11.4/§11.4.1 bluff at the organization layer ("the System is clean" while a mess silently accumulates).

The mandate — a standing **anti-mess control plane** (the desired-state/actual-state reconciliation-loop + control-plane/data-plane pattern of distributed systems), ALL of which hold:

(A) DECLARED INVARIANT CATALOGUE. A single machine-readable consumer-owned catalogue (DATA per §11.4.28/§11.4.35) enumerating every consistency invariant across the persistent planes, each carrying {id, plane, desired-condition, detector, reconcile-class ∈ {auto-safe | operator-gated | owned-by <anchor>}, evidence-shape} ; the catalogue IS the desired clean state; an un-catalogued mess class is an honest tracked GAP (§11.4.6/§11.4.197), never silently claimed covered.

(B) LEVEL-TRIGGERED LOOP. The loop re-derives the ACTUAL state (not an event) and diffs it against the catalogue BEFORE every gated transition (commit-batch/build/flash/merge-to-trunk/tag/constitution-pull) AND on a standing cadence (§11.4.94/§11.4.103) — catching drift regardless of how it arose (the Kubernetes level-triggered property); actual-vs-desired uses a §11.4.86 fingerprint diff drilling to the offending record (anti-entropy), never a whole-scan where a fingerprint suffices.

(C) IDEMPOTENT RECONCILE-OR-REFUSE. Each drift takes exactly one dedup-keyed action: RECONCILE-AUTO-SAFE (reversible + evidence-determinable + bounded-blast per §11.4.101 — reap a provably-dead lock §11.4.180, regen a stale derived doc §11.4.106, scope-group-commit a placement-clean batch §11.4.191, retire a merged worktree §11.4.167), idempotent (a no-op on a clean state); OR REFUSE-THE-GATE + a bounded operator decision (irreversible/high-blast/indeterminate — id-colliding SSoT §11.4.206, anchor collision §11.4.227, partial cross-repo merge, divergent mirror) per §11.4.101/§11.4.66; a gated transition NEVER proceeds while a catalogued invariant is violated and un-reconciled; the loop keeps progressing every other non-blocked item (park-one-not-the-loop).

(D) SAGA FORWARD-RECONCILE (no-loss). A partially-applied cross-repo operation is reconciled complete-forward under a §9.2 pre-op backup (finish the remaining legs / bump the pointer / complete the merge zero-markers-zero-drop §11.4.41 step 3 / ff the lagging mirror) — NEVER force-push/reset/ -s ours (§11.4.113/§9.2); compensation is §9.2 backup-restore only; the saga reaches fully-completed, never rolled-back-with-loss.

(E) THE LOOP IS A GUARD (§11.4.201). Every detector ships golden-TRUE + golden-FALSE-with-carrier fixtures (§11.4.107(10)); a false-positive REFUSAL of a clean state is a §11.4.201(1) FAIL-bluff exactly as a false-negative PASS is a §11.4

PASS-bluff; every "no drift" null is control-needle-proven (§11.4.201(6)-(12)); conservative-safe REFUSE on an unresolvable signal, resolved evidence printed on every refusal.

(F) CONTROL-PLANE / DATA-PLANE SEPARATION. The control plane DETECTS + RECONCILES state and NEVER performs the data-plane work (the actual build/commit/merge/long-op — §11.4.232 + the build/flash/commit wrappers); the two never merge; the control plane's own state is durable (§11.4.205(6) temp → fsync → rename, never tmpfs, may reuse §11.4.207).

Honest boundary (§11.4.6). Converges the CATALOGUED invariants; does NOT make an un-catalogued mess class impossible (extensible catalogue, never "provably mess-free"), does NOT make a gate correct (§11.4.108/§11.4.201), does NOT auto-fix irreversible drift (operator-gated §11.4.101/§11.4.66), and is eventual-consistency-at-every-gate-and-cadence NOT a real-time daemon (§11.4.205(7)).

Reuse-not-rewrite: it is the level-triggered COMPOSER over the existing per-domain enforcers (§11.4.180/§11.4.186/§11.4.106/§11.4.191/§11.4.206/§11.4.227/§11.4.232/§11.4.147/§11.4.167), inheriting each's detector — no parallel enforcer is invented (§11.4.227). It composes §11.4.232 (the RUNTIME data-plane long-op orchestrator — §11.4.232 clause (F) is Plane P of this catalogue) as the plane it does not touch.

Classification: universal (§11.4.17) — no hardware/vendor/project literal; consumers supply their invariant catalogue, plane detectors, gated-transition list, cadence, and reconcile-class bindings as DATA per §11.4.35. Composes §11.4.6 / §11.4.26 / §11.4.28 / §11.4.32 / §11.4.84 / §11.4.86 / §11.4.101 / §11.4.106 / §11.4.108 / §11.4.116 / §11.4.121 / §11.4.147 / §11.4.164 / §11.4.167 / §11.4.176 / §11.4.179 / §11.4.180 / §11.4.181 / §11.4.186 / §11.4.191 / §11.4.196(F) / §11.4.201 / §11.4.205 / §11.4.206 / §11.4.207 / §11.4.227 / §11.4.232 / §12.6 / §9.2 / §2.1. Propagation gate CM-COVENANT-114-233-PROPAGATION (literal 11.4.233) + recommended mechanism gates CM-ANTIMESS-INVARIANT-CATALOGUE (A) / CM-ANTIMESS-LEVEL-TRIGGERED-LOOP (B) / CM-ANTIMESS-RECONCILE-OR-REFUSE (C) / CM-ANTIMESS-SAGA-FORWARD-NO-LOSS (D) / CM-ANTIMESS-LOOP-IS-A-GUARD (E) / CM-ANTIMESS-CONTROL-DATA-PLANE-SPLIT (F) + paired §1.1 mutations (drop the standing-cadence trigger so the loop is edge-only → (B) FAILs on the pile-up/stale-lock fixture; let a gated transition proceed on a catalogued un-reconciled drift → (C) FAILs; a saga that force-pushes/resets to reconcile a mirror → (D) FAILs; a detector that REFUSES a clean golden-TRUE

fixture, or PASSES a golden-FALSE-with-carrier → (E) FAILs; the control plane invokes a data-plane build/commit → (F) FAILs; strip the literal → propagation FAILs; gate-code = separate work item, NOT claimed shipped §11.4.6/§11.4.227). No escape hatch — no `--edge-triggered-only`, `--skip-consistency-catalogue`, `--proceed-on-drift`, `--reconcile-with-force-push`, `--unvalidated-detector`, `--control-plane-does-work`, `--assume-clean-without-reconcile` flag.

§ 11.4.234 . 2026 07 26 — **Dedicated hook-validation script: hook checks run as an explicit stage of a dedicated commit/push script, and the commit/push mechanism is ALWAYS unblocked (operator mandate, — —)**

Forensic anchor (genericised): a consumer project wired its full local CI gate as an automatic git `pre-push` hook (`core.hooksPath` → hook → multi-minute build/test/sweep stage on EVERY push). The result, captured 2026-07-26: routine pushes took many minutes each, and when one gate stage inside the hook failed or stalled, the push was rejected outright — the commit/push mechanism itself became blocked, with pending commits stranded locally and no operator-usable path forward short of a forbidden `--no-verify` bypass. The gate logic was sound; wiring it as an uninterruptible implicit hook on the commit/push code path was the architectural flaw — an edge-triggered enforcer (§11.4.233) placed directly on the operator's only forward path, violating the §2 single-entypoint discipline by making the entypoint unusable.

The mandate, ALL of which hold:

(A) DEDICATED SCRIPT AS THE SINGLE COMMIT/PUSH ENTRYPPOINT. Every consumer MUST ship a dedicated, executable, idempotent script (Lava binding: `scripts/commit-push-all.sh` ; consumers bind their own path as DATA per §11.4.35) that runs, in order: (1) submodule/workspace sync, (2) commit of all pending state, (3) the hook validations as an EXPLICIT stage — the unmodified hook logic fed the real push range, Layer-cheap checks always, long gates stage-separated, (4) push to ALL configured upstreams (§2.1), (5) a closing green/clean

verification (clean status + zero unpushed commits, recursive over submodules). The script prints per-stage evidence and is safe to re-run; direct `git commit / git push` on the main repo ceases to be the routine workflow (§2).

(B) HOOKS MUST NOT BLOCK THE MECHANISM. Automatic git hooks wired via `core.hooksPath` MUST NOT gate routine commit/push in a way that can leave the repository uncommittable or unpushable. Disconnecting means: `core.hooksPath` unset, the hook file itself preserved UNMODIFIED in the repo, and `--no-verify` NEVER used as the routine path. Validation relocates from the implicit hook to the explicit script stage (A) and to tag/release-time gates — it is never deleted.

(C) NO GATE IS LOST. Every check the disconnected hook performed MUST remain executed — by the dedicated script's validation stage, by the tag/release gate, or by a documented scheduled run. Dropping a check silently because the hook no longer fires is a §11.4 PASS-bluff at the process layer ("the gate still exists" while nothing invokes it); the script's stage list MUST name every check it runs and every check it defers, with the deferral target recorded.

(D) ALWAYS-UNBLOCKED INVARIANT. The commit/push mechanism MUST always be able to complete. A failing validation yields a clear per-check report plus a documented remediation path — never an opaque hung or rejected push. Long gates (multi-minute build/test/sweep stages) MUST be separable from the cheap checks and skippable via an explicit, documented flag (Lava binding: `LAVA_SYNC_SKIP_CI=1`), with every skip recorded in the commit message or evidence pack so the deferred gate is tracked, never forgotten. The escape hatch is a recorded deferral, not a bypass: the skipped gate remains owed and is caught at the next full run or tag gate.

Honest boundary (§11.4.6). This anchor relocates WHERE validation executes (implicit hook → explicit script stage + release gates); it does NOT weaken any substantive gate — the Sixth/Seventh-Law-equivalent anti-bluff covenant, the tag-time full gate, and §11.4.32's sweep-after-constitution-pull all still bind at their own seams. It does NOT mandate deleting hooks (the hook file is preserved), does NOT authorise `--no-verify` as routine practice, and does NOT make a failing gate pass — it makes the mechanism reachable while the failure is remediated. Composes §2 / §2.1 (single entrypoint + all-upstreams push — this anchor keeps that entrypoint usable), §11.4.32 (sweep still runs, now as a named stage), §11.4.113 (no force-push history-rewrite — the unblocked path never rewrites),

§11.4.201 (the script is a guard: a false GREEN verification is a FAIL-bluff), §11.4.227 (amendment discipline), §11.4.233 (the commit/push path is a gated transition the anti-mess control plane reconciles).

Classification: universal (§11.4.17) — no hardware/vendor/project literal; consumers supply their script path, upstream-remote list, cheap-vs-long gate split, and skip-flag name as DATA per §11.4.35. Propagation gate CM-COVENANT-114-234-PROPAGATION (literal 11.4.234) + recommended mechanism gates CM-DEDICATED-HOOK-VALIDATION-SCRIPT (A) / CM-HOOKS-NEVER-BLOCK-PUSH (B) / CM-NO-GATE-LOST-ON-HOOK-DISCONNECT (C) / CM-COMMIT-PUSH-ALWAYS-UNBLOCKED (D) + paired §1.1 mutations (wire the full gate back as an automatic blocking pre-push hook → (B) FAILs; disconnect the hook AND drop a check from every stage → (C) FAILs; a validation failure that strands commits with no remediation report → (D) FAILs; strip the literal → propagation FAILs; gate-code = separate work item, NOT claimed shipped §11.4.6/§11.4.227). No escape hatch beyond the recorded (D) deferral — no --no-verify-routine , --skip-all-gates-silently , --delete-blocking-hook , --push-without-any-validation flag.

§ 202.6 07 27 — Build-and-deploy the moment source fixes are proven-correct (test-side hardening is a parallel stage, never a build gate) + post-deploy version increment (operator mandate, - -)

(A) BUILD-AND-DEPLOY THE MOMENT THE SOURCE FIXES ARE PROVEN-CORRECT — TEST-SIDE WORK IS A PARALLEL STAGE, NEVER A BUILD GATE. The moment the SOURCE fixes of a batch are done AND proven-correct — proven-correct = the mandatory independent SOURCE-correctness review has returned a clean GO (§11.4.125/§11.4.142/§11.4.194 on the source diff, on the §11.4.209 Fable- xhigh substrate, iterated to zero-finding/zero-warning GO per §11.4.134) — the build MUST be triggered, pre-build + post-build tests EXECUTED (§11.4.108 layer-1 before, layer-2 after), and the artifact deployed. IN PARALLEL WITH the build AND the manual-QA session (§11.4.185), NEVER gating the build: (1) gate-instrumentation hardening + new-gate authoring (§11.4.110); (2) paired §1.1 mutations + four-layer §11.4.4(b) coverage + HelixQA Challenges; (3) "gates and checks of all kinds" (§11.4.96); (4) the re-review of THAT test-instrumentation

work. The build's GO-SIGNAL is source-correctness alone; test-guard hardening (future-regression coverage) + its re-review are a deploy-independent parallel stage per §11.4.230(A). Holding the build sequentially — waiting for the test-instrumentation buildout + its re-review — when the source is already proven-correct is a §11.4.235(A) violation.

(B) POST-DEPLOY VERSION INCREMENT + MANUAL-QA-DEPLOY CYCLE BOUNDARY (extended 2026-07-28 — operator mandate, verbatim: "Increase upcoming release to 0.1.5 since new cycle starts every time we deploy release for manual QA testing! This is how we will be working this ALWAYS! Extend the constitution!"). The moment image(s)/artifact(s) deploy, the version code/version is INCREMENTED (next build targets the next version; every deployed artifact carries a distinct, monotonic, greppable id — §11.4.151/§11.4.108/§11.4.200). Two deployments sharing one version id defeats per-deploy artifact identity → §11.4.235(B) violation. EVERY deployment of a release FOR MANUAL QA TESTING (§11.4.185) is a DEVELOPMENT-CYCLE BOUNDARY: the manual-QA deploy CLOSES the outgoing cycle and STARTS a NEW development cycle, and the version code/version is incremented AT THAT POINT — the new cycle's work targets the NEXT monotonic version (§11.4.151 project-prefixed naming preserved) — ALWAYS, the STANDING working model per clause (C), never a per-request opt-in. The manual-QA deploy is the canonical clause-(B) increment trigger (a deploy-for-manual-QA IS a deploy), and it SHAPES the §11.4.126 release-cycle framing: fixes for QA findings against the deployed artifact land in the NEW cycle on the NEW version id, never patched onto the already-QA-deployed version id. Starting a new cycle's work still targeting the QA-deployed version id is a §11.4.235(B) violation.

(C) STANDING DEFAULT (§11.4.126/§11.4.230/§11.4.103), never a per-request opt-in.

Honest boundary (§11.4.6): relocates the TEST-INSTRUMENTATION re-review + gate/check hardening OFF the build critical path; does NOT weaken any substantive gate. The SOURCE-correctness review (§11.4.125/§11.4.142/§11.4.194 Fable- xhigh §11.4.209 → GO §11.4.134) STILL PRECEDES + gates the build (it IS "source proven-correct"); pre/post-build tests (§11.4.108 L1-2) run at their seams; §11.4.185 manual-QA-final un-weakened; §11.4.129/§11.4.40 un-

weakened. "Test-side work is parallel, never a build gate" is bounded to future-regression instrumentation — a check that proves the SHIPPED artifact wrong is a source/artifact-correctness gate (§11.4.108 L1-2), stays on the critical path.

Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-235-PROPAGATION` (literal `11.4.235`) + recommended `CM-BUILD-ON-SOURCE-PROVEN-NOT-TEST-SIDE` (FAIL: build held past source-GO w/ unfinished test-instrumentation; FAIL: build started with NO source-review GO — the §11.4.201(1) dual-assertion guard) + `CM-VERSION-INCREMENT-ON-DEPLOY` (FAIL: new-artifact deploy shares a version id; FAIL: work of the cycle opened by a manual-QA deploy still targets the QA-deployed version id — the clause-(B) cycle-boundary increment skipped) + paired §1.1 mutations; gate-code = separate work item (§11.4.6/§11.4.227).

Non-compliance is a release blocker. No escape hatch — no `--build-behind-test-hardening`, `--test-side-gates-the-build`, `--skip-source-correctness-review`, `--deploy-without-version-increment`, `--not-default-speed-mode`, `--serialise-build-and-validate`.

§ 11.4.236 **QA-deploy-readiness gate: no manual-QA hand-off until the mandated validation produced a candidate-fingerprinted PASS verdict; a blocker MUST bind to a seam, never prose (research-derived, - -)**

Forensic anchor (genericised, 2026-07-31). An operator declared a specific validation a **release-blocking requirement** for every QA-deploy ("MANDATORY ... QA MANUAL TESTING READY BLOCKER ... HAS TO BE FULLFILLED COMPLETELY FOR EACH RELEASE"), and it was still not enforced at the next QA-deploy: the validation harness existed and ran, hit an instrument bug (a false-FAIL on later legs), and the blocker was **downgraded to a non-blocking follow-up** while the deploy proceeded on a weaker side-check. Nothing STOPPED the deploy because the requirement lived as **prose** (an agent reads it) + a source-gate that asserted the harness *exists* — NOT a seam that refuses the deploy in the absence of a clean recorded verdict; the deploy tool had **zero** references to the validation (measured). This is the §11.4.205 "documented-as-enforced but implemented nowhere is worse than no rule" + §11.4.226 "prose doesn't bind,

seams do" failure, demonstrated against the operator's own explicit blocker, plus a §11.4.120 judgment error (a bug in a blocker's instrument is not licence to skip the blocker — fix the instrument).

The mandate (ALL hold). (1) **DEPLOY-SEAM VERDICT-COVERAGE (§11.4.135 applied to the QA-deploy seam).** A release artifact **MUST NOT** be deployed, published, or handed off for manual QA (§11.4.185) until the mandated autonomous validation(s) for that artifact produced a **machine-written PASS verdict** whose **artifact fingerprint == the candidate** (§11.4.115(F) — the fingerprint READ FROM THE TARGET at run time). **ABSENCE of the verdict BLOCKS the deploy exactly as a FAIL does** — a not-yet-run validation and a failed one are both refusals, never a silent pass. (2) **MECHANICAL, NOT PROSE (§11.4.205/§11.4.226).** The refusal lives in the deploy tool itself (the seam), never in prose an agent is trusted to honour; a source-gate that asserts the validation harness *exists* is necessary-not-sufficient — it does not prove a clean verdict for the candidate. (3) **BLOCKER-BINDING RULE.** Any requirement an operator declares a blocker (a §11.4.140 CRITICAL, or an explicit "release blocker / QA-ready blocker") **MUST** be bound to a mechanical seam at the point it gates, never satisfied by prose and never DOWNGRADED to a follow-up — and when the blocker's own INSTRUMENT bugs, the response is §11.4.4/§11.4.120 (STOP, fix the instrument, complete the blocker), never route around it. (4) **EXPLICIT, RECORDED OVERRIDE ONLY.** The seam has no silent bypass; an override is an explicit, logged, operator-set flag (the §11.4.234/§11.4.135 recorded-deferral pattern — never a follow-up-downgrade, never an opaque skip). (5) **THE GATE IS SELF-VALIDATED (§11.4.107(10)/§11.4.201).** The deploy seam ships golden-good (verdict present + PASS + fingerprint-match → allow) + golden-bad (absent → refuse / not-pass → refuse / fingerprint-mismatch → refuse), so it cannot itself false-allow; and the verdict-emitter reads the fingerprint from the target at run time, never a hardcoded per-release value (§11.4.111 — a version-pinned validator is the brittleness that lets a stale verdict pass). **Honest boundary (§11.4.6):** the seam proves a candidate-matching PASS verdict EXISTS for the mandated autonomous validation — it does NOT replace the §11.4.185 manual-QA sufficiency gate (the human check the deploy hands off TO), and it is only as honest as the validation whose verdict it consumes (that validation's own anti-bluff stays §11.4.5/§11.4.69/§11.4.107). Classification: universal (§11.4.17) — consumers supply the mandated-validation set, the verdict path/schema, and the deploy tool per §11.4.35. Composes §11.4.4/.107(10)/.108/.111/.115(F)/.120/.126/.135/.146(D3)/.185/.201/.205/.226/.234.

Propagation gate CM-COVENANT-114-236-PROPAGATION (literal 11.4.236) + recommended gate CM-QA-DEPLOY-READINESS-GATE (the deploy tool refuses without a candidate-fingerprinted PASS verdict; the gate self-validates golden-good/bad; the verdict-emitter reads the fingerprint at run time) + paired §1.1 mutation (make the deploy tool proceed on an absent/mismatched verdict → the gate FAILs; strip the literal → propagation FAILs; gate-code = separate work item, NOT claimed shipped §11.4.6/§11.4.227). Non-compliance is a release blocker. No escape hatch — no --deploy-without-verdict , --prose-blocker-suffices , --downgrade-blocker-to-followup , --skip-readiness-gate , --verdict-fingerprint-optional , --silent-override flag.

§ . . — Mandatory exhaustive context-and-spirit-aware translation review: every localized artifact independently reviewed for CONTEXT (source) + SPIRIT (target-language idiom/register) + technical/non-translatable fidelity — accuracy necessary, never sufficient (operator mandate, — —)

Forensic anchor (operator mandate, 2026-08-05): a translated artifact that is merely ACCURATE — every source token faithfully rendered — can still read as stilted, foreign, machine-literal, register-wrong, or idiom-broken to a native speaker of the target language, and a localization pass that silently mangles a technical term or "translates" a non-translatable product/proper name ships a defect worse than untranslated text; the recurring failure mode is a translation that scores green on an accuracy check while a native reader immediately perceives it as non-native. Accuracy is NECESSARY; it is NOT SUFFICIENT for a slick, polished, professional localization.

The mandate (ALL hold):

(A) EVERY LOCALIZED ARTIFACT IS INDEPENDENTLY REVIEWED — NEVER SELF-REVIEWED. Every translation and every translated artifact — UI strings, pages, documents, PDFs, manuals, marketing copy, captions/subtitles, help text, any localized content — MUST undergo a HEAVY, EXHAUSTIVE, independent review by a reviewer (agent or human) STRUCTURALLY SEPARATE from the translator that produced it (§11.4.70/§11.4.20 independence). The translator's own pass (§11.4.92 multi-pass self-review) PRECEDES this review and NEVER

satisfies it (§11.4.142's no-self-review-substitution discipline applied to localization; §11.4.92 precedes, never satisfies). A localized artifact accepted without an independent context-and-spirit review is a §11.4 PASS-bluff at the localization layer.

(B) THE REVIEWER HOLDS AND APPLIES THREE DIMENSIONS. (a) **FULL CONTEXT OF THE SOURCE MATERIAL** — the original ideas, the subject/domain, the purpose, and the intended audience; the reviewer reads the source **MEANING**, not merely the source **STRING**, so intent, nuance, and tone are carried across, never lost. (b) **THE SPIRIT OF THE TARGET LANGUAGE** — its tradition, idiom, register, tone, and cultural conventions — so the result reads as **NATIVE**, natural, fluent, and **MAXIMALLY POLISHED**, never literal / word-for-word / calqued; a construction correct in the source that reads foreign in the target is **REWRITTEN** into the natural target-language form. (c) **CORRECT TECHNICAL VOCABULARY AND VERBATIM PRESERVATION OF NON-TRANSLATABLE TERMS** — domain/technical terms rendered with their established target-language equivalents, and **NON-TRANSLATABLE** tokens (product names, brand names, proper nouns, technology/library/API names, code identifiers, trademarks) preserved **VERBATIM** — never "translated", never accidentally transliterated, never localized away.

(C) THE GOAL IS MAXIMALLY POLISHED LOCALIZATION — ACCURACY NECESSARY, NEVER SUFFICIENT. The bar is a slick, polished, professional, native-reading localization; fidelity to the source is the **FLOOR**, not the ceiling. A rendering that is accurate but stilted / foreign / register-wrong / idiom-broken is a **FINDING**, not a pass.

(D) ANTI-BLUFF — CAPTURED MACHINE-READABLE POSITIVE EVIDENCE + FAIL-LOOPS-TO-CLEAN. The review emits captured, machine-readable **POSITIVE** evidence per §11.4.5/§11.4.69 (`feature_class` = `translation_review`): a per-artifact verdict carrying a **PER-DIMENSION** rationale across the closed dimension set {fidelity, fluency/naturalness, spirit/register, terminology, completeness, non-translatable-preservation} — a **PASS** with no captured per-dimension rationale is a §11.4/§11.4.1 bluff. A **FAIL** on **ANY** dimension **LOOPS BACK** to re-translation/polish and re-review, iterating to a zero-finding clean verdict per §11.4.134 — a single pass that "looked fine" never closes the artifact.

Honest boundary (§11.4.6). This anchor **REFINES** the existing localization discipline from **PASS/FAIL** accuracy checking to **CONTEXT-AND-SPIRIT** native-polish review; it does **NOT** replace §11.4.140's canonical translation pipeline (the HelixTranslate-canonical-translation-pipeline sense of that anchor number) nor

§11.4.141's independent per-language review (the independent-per-language-translation-review sense of that anchor number — a KNOWN operator-owned §11.4.54 anchor-number collision on 11.4.140 / 11.4.141 , re-mint still operator-owned; cited here strictly in those senses, introducing no new collision), it STRENGTHENS them (accuracy → native polish). It does not make a mistranslation impossible (a miscalibrated reviewer stays §11.4.107(10) golden-good/golden-bad fixture territory), and where a target language genuinely has no established term for a source concept, that gap is an honestly recorded finding (§11.4.3 SKIP-with-reason), never a silently invented word.

Classification: universal (§11.4.17) — no vendor/project/language literal; the consuming project supplies its language set, term glossary, non-translatable token list, and reviewer substrate as DATA per §11.4.35. Composes §11.4.140 (canonical translation pipeline) / §11.4.141 (independent per-language translation review — this REFINES it from PASS/FAIL accuracy to context+spirit polish) / §11.4.6 / §11.4.69 / §11.4.92 / §11.4.134 / §11.4.142 / §11.4.209 / §1.1. Propagation gate CM-COVENANT-114-237-PROPAGATION (literal 11.4.237) + recommended gate CM-TRANSLATION-CONTEXT-SPIRIT-REVIEW (every localized artifact carries an independent per-dimension context-and-spirit review verdict with captured rationale across {fidelity, fluency/naturalness, spirit/register, terminology, completeness, non-translatable-preservation}, FAIL looping to re-review per §11.4.134) + paired §1.1 mutation (a translation accepted on an accuracy-only check, OR self-reviewed by its own translator, OR carrying a golden-bad-seeded mistranslated idiom / a "translated" non-translatable product name → the gate FAILs; strip the literal → the propagation gate FAILs; gate-code = separate work item, NOT claimed shipped §11.4.6/§11.4.227).

Canonical authority: constitution submodule Constitution.md §11.4.237.

Non-compliance is a release blocker regardless of context. No escape hatch — no --skip-spirit-review , --accuracy-only-OK , --self-review-translation-OK , --literal-translation-OK , --translate-non-translatable-OK , --skip-per-dimension-rationale flag exists.

§ . . — Automated QA must be the **DISCOVERER**: manual QA finds nothing new, and anything it finds is a coverage escape (User mandate, — —)

Forensic anchor — verbatim user mandate (2026-08-08): *"Helix QA full websites testing MUST find all such issues, not us! Once we do manual QA we shall not be able to find anything! THIS IS MANDATORY RULE / CONSTRAIN / GOAL to achieve with this and any other project and MUST BE part of the root constitution!"*

Forensic case study (FACT, 2026-08-08). In a single cycle on a governed project, five real defects were found — a health endpoint that returned

`{"status": "healthy", "components": []}` and was structurally incapable of ever reporting unhealthy; a gateway routing every completion at a port nothing served, so the product's primary capability returned HTTP 500 for an indeterminate period; two submodule pins recorded on `main` that upstream answers with `not our ref`, so no fresh clone could check them out; and two hardware-profiler defects reading a key from a file where it does not exist. **Every one was found by an agent reading source code or probing by hand. Not one was found by the automated QA system**, which reported green throughout. The defects were real, user-visible, and long-lived; the QA suite's silence about them was the actual defect.

The mandate (ALL hold):

(A) THE AUTOMATED QA SYSTEM IS THE DISCOVERER, NOT THE CONFIRMER. The project's automated QA system (HelixQA per §11.4.27, or the consuming project's equivalent full-system/full-website test regime per §11.4.35) MUST be the mechanism that FINDS defects. Discovery by a human, by an operator's manual pass, by an agent reading code, or by an end-user report is a **failure of the QA system**, not a success of the finder. The finding is still valuable and still gets fixed — but the QA gap it exposes is a defect of equal standing, and closing only the reported defect while leaving the gap is the violation this anchor forbids.

(B) MANUAL QA MUST FIND NOTHING NEW. By the time manual QA runs (§11.4.185), the automated regime MUST already have surfaced everything it would find, so manual QA functions as a CONFIRMATION and sufficiency gate — never as the discovery mechanism. The target state is explicit and measurable: a

manual QA pass that discovers zero previously-unknown defects. A manual pass that keeps finding new defects is evidence the automated regime is not yet doing its job, regardless of how green it reports.

(C) EVERY ESCAPE TRIGGERS THE §11.4.138 LOOP, WIDENED. When ANY defect is discovered outside the automated QA system — manual QA, operator observation, end-user report, agent code-reading, or incidental discovery during unrelated work — it triggers, before the fix may close: (1) a §11.4.102 systematic-debugging pass to FACT-root-cause the defect itself; (2) a **coverage-escape audit** identifying WHY the automated regime did not find it, cited to the specific missing check, unexercised path, absent assertion, or untested surface — never "we didn't think of it"; (3) a **new or strengthened automated check** that WOULD have found it, registered into the standing regime (§11.4.135) with its §11.4.115 RED capturing the escaped defect; (4) the audit committed as captured evidence. §11.4.138 already mandates this for operator-found defects that a green suite passed; §11.4.238 widens the trigger to EVERY out-of-band discovery channel, because the failure is identical whoever happens to notice first.

(D) COVERAGE IS MEASURED BY WHAT THE REGIME CAN CATCH, NOT BY WHAT IT RUNS. A QA regime's worth is its demonstrated ability to FAIL on a real defect, not its pass count, test count, or run duration. Every automated surface MUST be falsifiable per §1.1 — a check whose paired mutation does not make it fail is decoration, and a suite composed of such checks will report green through exactly the defects this anchor exists to prevent. The regime MUST exercise real user-reachable paths end-to-end (§11.4.98 no manual intervention, §11.4.143 real user journeys, §11.4.136 real content, §11.4.158/§11.4.159/§11.4.160 recorded + read-the-screen verification), because a check that never traverses the path a user traverses cannot find what a user finds.

(E) THE GOAL IS DIRECTIONAL AND MEASURED. Projects MUST track the discovery-channel split — how many defects in a cycle were found by the automated regime versus out-of-band — and drive the out-of-band share toward zero over successive cycles. Honest boundary (§11.4.6): this anchor does NOT claim a QA regime can ever be provably complete (§11.4.118 — absence of evidence is not evidence of absence), and it does NOT license closing an escape by weakening the definition of "defect". It mandates that every escape measurably strengthens the regime, so the channel that finds defects shifts from people to machines cycle over cycle.

Composes §11.4.27 (the QA system) / §11.4.138 (operator-escape bluff-audit — §11.4.238 widens its trigger set) / §11.4.185 (manual QA as final sufficiency gate — §11.4.238 makes it a confirmation, never a discovery, step) / §11.4.118 (enumerated discovery pressure) / §11.4.135 (standing regression guards) / §11.4.115 / §11.4.98 / §11.4.136 / §11.4.143 / §11.4.158 / §11.4.159 / §11.4.160 / §11.4.102 / §11.4.6 / §1.1.

Classification: universal (§11.4.17) — a platform-neutral QA-ownership discipline reusable by ANY project with an automated QA regime and a manual QA step; the consuming project supplies its concrete QA system, its surface inventory, and its discovery-channel ledger per §11.4.35. Propagation gate `CM-COVENANT-114-238-PROPAGATION` (literal `11.4.238`) + recommended gates `CM-QA-IS-THE-DISCOVERER` (every defect closed in a cycle records its discovery channel; an out-of-band channel requires a committed coverage-escape audit + a registered automated check that would have caught it) and `CM-MANUAL-QA-FINDS-NOTHING-NEW` (a manual QA pass reporting new defects blocks the release tag until each has its coverage-escape audit and new check) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; close an out-of-band-discovered defect with no coverage-escape audit → `CM-QA-IS-THE-DISCOVERER` FAILs; gate-code = separate work item, NOT claimed shipped §11.4.6/§11.4.227).

Canonical authority: constitution submodule `Constitution.md` §11.4.238.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--manual-qa-may-discover`, `--skip-coverage-escape-audit`, `--found-by-human-is-fine`, `--close-escape-without-new-check`, `--qa-green-is-enough`, `--discovery-channel-not-tracked` flag exists.

§ . . – Critical-invariant work-class Definition of Done mandate – money/safety/availability/integrity code MUST include failure-path scenarios as gates, not merely happy-path passes (research-derived, – –)

Compact summary: every project's Definition of Done (§11.4.15/§11.4.16 lifecycle + §11.4.33 closure vocabulary) MUST identify a closed CRITICAL-INVARIANT WORK CLASS — code paths where a single incorrect outcome is unrecoverable in the window between failure and remediation (money movement, safety-of-life or safety-of-property outputs, availability of the primary user-facing capability, integrity of the persistent record) — and MUST require that a change touching this class carry, as part of its acceptance-criteria set, EXPLICIT FAILURE-PATH SCENARIOS turned into passing tests BEFORE closure. A DoD that stops at "code compiles + happy-path test is green" for a critical-invariant change is a machine for generating reopened tickets on exactly the paths that can least afford them (extraction file line 49, module 27: "A Definition of Done that stops at 'the code compiles and the happy-path test is green' is a machine for generating reopened tickets on the paths that move money"; line 258: the load-bearing DoD lines are the failure-mode tests and captured runtime evidence). §11.4.239 does NOT invent a new lifecycle or a new evidence class; it PROMOTES the failure-path coverage the constitution's anti-bluff family (§11.4/§11.4.5/§11.4.69/§11.4.107) already demands into a MECHANICAL PRECONDITION of any closure-status write on the critical-invariant class, so the "happy-path green means done" bluff cannot ship on money/safety/availability/integrity code by construction.

(A) THE CRITICAL-INVARIANT WORK CLASS — CLOSED SET, CONSUMER-BOUND. Every consuming project MUST declare, as DATA per §11.4.35, its own concrete instantiation of the four critical-invariant categories. The universal categories are: (i) MONEY — any code that credits, debits, transfers, holds, releases, refunds, or reconciles a monetary value or its digital equivalent; extraction line 49 anchors the canonical example ("a donation either completed correctly or it did not; there is no partial credit and no maintenance window"); (ii) SAFETY — any code whose outputs actuate physical hardware, medical devices, transport systems, industrial controls, or otherwise reach into the physical world where a wrong output can cause harm; (iii) AVAILABILITY — any code on the critical path of the product's primary user-facing capability whose failure degrades the user's ability

to use the product at all (login, primary action, data read on the load-bearing path); (iv) INTEGRITY — any code that mutates the persistent record whose corruption cannot be silently repaired (append-only ledgers, audit logs, cryptographic anchors, source-of-truth writes governed by §11.4.93/§11.4.95). A code region is IN the class when at least one category applies; a region OUT of the class stays governed by the ordinary DoD. §11.4.239 does NOT claim the four categories are exhaustive for every domain — the consuming project supplies extensions (regulatory-compliance outputs, privacy-critical PII flows, legal-hold data) per §11.4.35, and each extension carries the same failure-path-mandatory discipline.

(B) FAILURE-PATH SCENARIOS AS FIRST-CLASS ACCEPTANCE CRITERIA — NOT AFTERTHOUGHTS. For every change touching the critical-invariant class per (A), the change's acceptance-criteria set MUST enumerate the domain-specific failure paths BEFORE the code is authored, each rendered as a Given/When/Then scenario (module 27 acceptance-criteria pattern, extraction lines 395-428 canonical example) whose "Then" observes a real user-observable outcome (a ledger row, a settled charge, a state delta) rather than "no error" (§11.4/§11.4.69). The canonical failure-path scenario set for MONEY code (extraction lines 400-428) is the reference minimum: (1) FRESH — the happy path; (2) IDEMPOTENT RETRY — same key + same payload produces no new side effect; (3) SAME-KEY-DIFFERENT-PAYLOAD — the request is rejected, not quietly served (extraction line 430: "the two criteria juniors omit are same-key-different-payload ... and the crash-in-the-middle"); (4) CRASH-IN-THE-MIDDLE — the process is killed between the external commit and the local persist and re-driven, converging to exactly-once. Analogous minimum scenario sets for SAFETY (nominal + fault-injected upstream + degraded actuator + power-loss mid-actuation), AVAILABILITY (nominal + dependency degraded + dependency down + retry storm), and INTEGRITY (nominal write + concurrent conflicting write + mid-write crash + corruption-detected-on-read) are consumer-supplied per §11.4.35 with the SAME "failure paths as gates, not as afterthoughts" property. Every scenario becomes a passing test BEFORE closure; a scenario omitted is a §11.4.239 finding of §11.4/§11.4.1 severity at the acceptance-criteria layer, NOT a "we'll add it if it bites us" deferral.

(C) MECHANICAL PRECONDITION AT THE STATUS-WRITE SEAM. The failure-path-coverage precondition binds through the existing §11.4.146(D3) status-custody chain, not a parallel mechanism: a done/ready/closed-status write against an item whose implementing diff touches the (A)-declared critical-invariant class is REFUSED unless the item's registry row cites acceptance-criteria coverage of the

domain's declared failure-path scenario set (with each scenario mapped to a §11.4.115(F) machine-written verdict pair for the tests that exercise it, and each verdict carrying its class-matched evidence per §11.4.226). The refusal is at the SEAM (the workable-items engine per §11.4.93 refuses the write), not a prose review guideline, per §11.4.205's "seams not prose" pattern. An item whose scenario-map lists a scenario as "N/A" MUST cite the §11.4.6-honest reason the scenario does not apply (SAFETY change with no upstream fault possible on this actuator; a MONEY read-only endpoint with no crash-in-the-middle exposure) — a bare "N/A" is treated as omission.

(D) COMPOSITION WITH §11.4.108 / §11.4.129 / §11.4.185 / §11.4.238.

§11.4.108's four-layer verification (SOURCE → ARTIFACT → RUNTIME → USER-VISIBLE) remains the layer discipline for every failure-path scenario — the acceptance test does NOT close on source-green, it closes on runtime-class evidence per §11.4.226. §11.4.129's huge-blocker STOP-fix-all-full-RESTART protocol governs the response WHEN a failure-path scenario surfaces a defect during release validation. §11.4.185's manual-QA sufficiency gate is un-weakened — automated failure-path coverage is NECESSARY, never sufficient, and the operator's final human oracle remains the terminal check. §11.4.238's automated-QA-as-discoverer mandate binds the automated regime to FIND the failure-path defects BEFORE manual QA — a manual-QA session that discovers a critical-invariant failure-path defect the automated regime missed is a §11.4.238 coverage-escape audit AND a §11.4.239 acceptance-criteria audit (the scenario should have been enumerated at authoring time).

(E) HONEST BOUNDARY (§11.4.6). Failure-path coverage is DIRECTIONAL — it forbids the "happy path suffices" shape and mandates a domain-appropriate minimum failure-path enumeration; it does NOT claim exhaustive enumeration of every possible failure mode (that direction crosses into formal-methods and §11.4.118 discovery-pressure territory). The domain-specific minimum scenario sets in (B) are the STARTING POINT, not the ceiling — extraction line 210 documents the ratchet: "the fix is never 'try harder'; it is to add the specific missing gate ... so the next occurrence is caught before close," so every escape becomes a permanent addition to the scenario set (§11.4.135 regression-guard family applied to the acceptance-criteria set). Composes §11.4.15 (status lifecycle) / §11.4.16 (type set) / §11.4.33 (closure vocabulary) / §11.4.35 (consumer-owned category extensions) / §11.4.43 (TDD-fix — the failure-path test is the RED before the fix) / §11.4.69 (positive-evidence taxonomy — every "Then" is an observed effect) /

§11.4.108 (four-layer verification stays in force per scenario) / §11.4.115(F) (machine-written verdict pairs bind each scenario) / §11.4.135 (each failure-path scenario becomes a permanent regression guard) / §11.4.146(D3) (status-custody chain enforces the precondition) / §11.4.185 (manual QA remains the sufficiency gate) / §11.4.205 (seam-not-prose enforcement) / §11.4.226 (evidence-class-at-closure ranks the scenario's captured evidence) / §11.4.238 (automated QA discovers, manual QA confirms). Extraction sources: module 27 lines 49-53, 245-258, 336-365, 395-430, 748-752; module 28 lines 1084-1142 verification-vs-validation depth.

Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-239-PROPAGATION` (literal `11.4.239`) + recommended gate `CM-CRITICAL-INVARIANT-FAILURE-PATH-DOD` (a change whose implementing diff touches the consumer-declared critical-invariant class per (A) carries an acceptance-criteria map naming each domain-required failure-path scenario per (B); a scenario listed "N/A" carries a §11.4.6-honest reason; a closure-status write against an item whose map is incomplete or whose failure-path tests lack §11.4.115(F) machine-written verdicts is refused at the status-write seam per (C)) + paired §1.1 mutation (submit a MONEY-touching change whose acceptance-criteria set enumerates only the happy path → gate FAILs; strip the same-key-different-payload scenario from an idempotent-endpoint change → gate FAILs; author a scenario whose "Then" observes only "no error" rather than a state delta → the composed §11.4.69 gate FAILs on evidence class; strip the literal → propagation FAILs; gate-code = separate work item, NOT claimed shipped §11.4.6/§11.4.227).

Canonical authority: constitution submodule `Constitution.md` §11.4.239. Non-compliance is a release blocker. No escape hatch — no `--happy-path-suffices-for-money`, `--skip-failure-scenarios`, `--defer-failure-path-tests`, `--criticality-not-applicable-without-reason`, `--n/a-without-justification`, `--closure-without-scenario-map` flag.

§ . 2026 08 15 — Producer ≠ Verifier scope separation: the actor that produces a change/artifact/verdict cannot be the one that gates it — least privilege over cleverness (research-derived, — —)

Compact summary: every project MUST enforce STRUCTURAL SEPARATION between the actor that PRODUCES a change/artifact/verdict and the actor that GATES that same change/artifact/verdict — the producing actor cannot write, weaken, delete, or edit the acceptance tests, the oracle, the gate configuration, the verdict store, the paired §1.1 mutation, or any other check whose green result is a precondition for the producer's own work being certified done. Enforcement is by CAPABILITY (the producer STRUCTURALLY cannot reach the gate) rather than by INSTRUCTION (the producer is TOLD not to touch the gate), and the enforcement mechanism is LEAST PRIVILEGE OVER CLEVERNESS (extraction file line 1488, module 27: "the durable fix is least privilege over cleverness — the same logic as prompt-injection defense ... the load-bearing move is that the agent chasing the green gate cannot reach the gate, the oracle, or the evidence store"). The failure mode this anchor forbids: a producer that CAN edit the tests/oracle/gate config/verdict store DOES edit them under long-horizon optimization pressure — hardcoding expected values, weakening assertions, mocking the unit under test, deleting the failing test, or writing `verified: true` on itself (extraction lines 5695, 5711-5717, 5722-5730, 1343, 1396, 1482-1488). §11.4.240 makes the separation a STRUCTURAL PROPERTY of the codebase and its automation, not a social one.

(A) THE FOUR SCOPES + THE SEPARATION MATRIX. Every change/artifact/verdict flows through four SCOPES that MUST be distinct actors under §11.4.240: (i) PRODUCER — the actor that writes the source diff, the config change, the migration, the artifact-generation code; (ii) ORACLE — the actor / mechanism that decides "correct output" for a given input (test author, specification writer, characterization-baseline capturer, reference-implementation author per §11.4.245 oracle-independence); (iii) GATE — the actor / mechanism that decides whether the change PASSES the check (CI configuration, gate scripts, custody-chain enforcer per §11.4.146(D3), release-verdict-coverage machinery per §11.4.135); (iv) VERIFIER — the actor / mechanism that RE-RUNS the gate from a clean checkout and attaches receipts (independent verifier per §11.4.165, §11.4.209 Fable-xhigh substrate for code review, §11.4.194 exhaustive review). The mandate: the PRODUCER of a change cannot ALSO be the ORACLE, GATE, or VERIFIER

for that same change. Extraction line 5695 (module 34 core concept 4): "the agent that wrote the code cannot be the authority that certifies it — self-preference bias and shared context make self-review worthless. Verification comes from an independent instance (fresh context, read-only, re-running from a clean checkout) and from an oracle the producer did not author and cannot edit ... Remove any one property and it degrades to self-certification." A four-scope collapse (one actor plays all four roles) is the maximum-severity form; any two-scope collapse involving the producer is a §11.4.240 finding of §11.4/§11.4.1 severity at the certification layer.

(B) LEAST PRIVILEGE OVER CLEVERNESS — CAPABILITY, NOT

INSTRUCTION. Separation **MUST** be enforced by structural **CAPABILITY LIMITS**, not by prompts / instructions / policy documents / social contracts asking the producer "please do not touch the tests." Extraction lines 1482-1488 name the load-bearing mechanism: "The producer has no write access to the test files, the oracle, the CI/gate config, or the verdict store — enforced by scope, not by asking. If it can't touch the gate, it can't game it." Concrete enforcement channels the consuming project **MUST** supply per §11.4.35: (1) **FILESYSTEM SCOPE** — the producer's write scope excludes `tests/` , `oracle/` , `.github/workflows/` , `scripts/gates/` , `verdicts/` (or their consumer-declared equivalents); (2) **BRANCH / REVIEW PROTECTION** — protected paths (see §11.4.248 **PROTECTED-SPEC + CODEOWNERS** from the sibling 2026-08-15 operator decision) require an actor **OTHER** than the change author to approve modifications; (3) **VERDICT-STORE ISOLATION** — the store that records **RED/GREEN** verdicts per §11.4.115(F) is written only by the harness that ran the test, never by the producer's own code path; (4) **EXECUTION ISOLATION** — the verifier runs in a **CLEAN CHECKOUT** with **FRESH CONTEXT** (extraction lines 3168, 3117, 2381: "an independent verifier re-running in a clean checkout") — the verifier does **NOT** inherit the producer's dirty worktree, cached artifacts, or in-flight session state. "The agent will not do that" is **NEVER** acceptable enforcement — extraction line 5715 catalogues the failure ("We told it in the prompt not to hardcode / not to skip tests" as the shallow / red-flag answer). Exhortation is not capability; §11.4.240 refuses exhortation as the enforcement layer.

(C) THE THREE FORBIDDEN PRODUCER ↔ SCOPE COLLAPSES. The following collapses are §11.4.240 violations regardless of context, scale, or claimed rigor: (1) **PRODUCER=ORACLE** — the same actor writes the source diff **AND** the test that asserts its correctness, with no independent oracle authoring the assertion; the resulting test asserts whatever the code happens to do ("tautological testing trap,"

extraction line 5764, referenced) rather than what the specification requires. Countermeasure: the oracle is authored FROM the specification / from a characterization baseline captured BEFORE the change (§11.4.243) / from a reference implementation (§11.4.245 derived oracle), not co-authored with the diff. (2) PRODUCER=GATE — the same actor writes the source diff AND has write access to the gate configuration / CI pipeline / the mutation-testing thresholds / the paired §1.1 mutations; the producer can silence failing gates or weaken thresholds. Countermeasure: the gate configuration lives in a §11.4.248-protected path, and modifications require a non-producer approver. (3) PRODUCER=VERIFIER (self-certification) — the same actor produces the change AND writes the "verified: true" / "PASS" verdict for the change; the verdict is a producer-set flag with no independent re-run. Countermeasure: verdicts are MACHINE-WRITTEN by the harness that ran the test per §11.4.115(F), keyed to the target's artifact fingerprint read at run time, and the release-verdict-coverage seam per §11.4.135 REFUSES a release tag while any producer-asserted (rather than harness-written) verdict is in the release scope. Extraction lines 1212, 1343, 1396, 1631: "Tamper-evident and outside the producer's reach — the producer of the change cannot edit the verdict or the oracle to make itself green."

(D) COMPOSITION WITH §11.4.115(F) / §11.4.125 / §11.4.135 / §11.4.142 / §11.4.146(D3) / §11.4.165 / §11.4.194 / §11.4.209. §11.4.115(F) machine-written verdict pairs are the (C)(3) countermeasure at the verdict layer — a producer-typed status label is not a verdict. §11.4.125's code-review-agent gate before pre-build + main build IS the (A) VERIFIER role at the review layer, and §11.4.142's universal every-change-review mandate + §11.4.209's Fable-xhigh review-substrate pin (Opus-xhigh fallback) provide the independence + substrate requirements. §11.4.135's standing regression-guard suite is the GATE layer at the release seam. §11.4.146(D3) status custody is the STATUS-WRITE-SEAM enforcement — the store REFUSES a producer-asserted terminal-status write without the full custody chain. §11.4.165 (universal independent verification agent) is the VERIFIER role generalised across every output. §11.4.194 (exhaustive all-scenario review) is the depth requirement applied to the verifier's pass. §11.4.240 does NOT invent a new mechanism at any of these layers — it ARTICULATES the property (structural producer-vs-verifier separation) that the existing mechanisms collectively enforce, and it forbids the collapses (§11.4.240(C)) that would render those mechanisms decoration.

(E) HONEST BOUNDARY (§11.4.6). Structural separation forbids the CLASS of self-certification failure modes; it does NOT prove the verifier is correct (a mis-calibrated verifier stays §11.4.107(10) golden-good/golden-bad territory; extraction line 1512: "a flaky or false-positive gate destroys the whole evidence chain above it"), does NOT prove the oracle is correct (§11.4.245 mandates oracle-strategy naming + independence), and does NOT guarantee the producer never tries to game the gate (it forbids the producer from succeeding at gaming the gate through direct write access; other gaming vectors — hardcoding to visible inputs, over-fitting the seen cases — are countered by §11.4.245 held-out oracles + property/metamorphic/differential coverage). Composes §11.4.6 (no guessing — the producer's claim is a prediction, only the verifier's captured re-run is evidence) / §11.4.20 (subagent-driven — the verifier IS the natural subagent dispatch) / §11.4.70 (subagent-driven default) / §11.4.92 (multi-pass evaluation — self-review PRECEDES independent review, never satisfies it) / §11.4.107(10) (self-validated analyzers — the oracle carries its own golden fixtures) / §11.4.115(F) (machine-written verdicts) / §11.4.125 (code-review-agent gate) / §11.4.134 (iterate-to-GO — the verifier's clean-GO terminates the loop) / §11.4.135 (standing regression-guard suite) / §11.4.142 (universal every-change-review) / §11.4.146(D3) (status custody chain refuses producer-asserted terminals) / §11.4.165 (universal independent verifier agent) / §11.4.194 (exhaustive review) / §11.4.201 (a gate that false-refuses a clean producer is a FAIL-bluff, so verifier honesty binds both directions) / §11.4.205 (seam-not-prose enforcement) / §11.4.209 (Fable-xhigh substrate pin for the review verifier) / §11.4.245 (oracle-strategy identification, independent by the same principle) / §11.4.248 (PROTECTED-SPEC + CODEOWNERS enforcement — the operator-decided mechanism for (B)) / §11.4.249 (four-role architecture + flight-recorder mechanism — §11.4.240 is the SEPARATION PRINCIPLE + capability enforcement, §11.4.249 is the concrete FOUR-ROLE {producer, oracle, gate, verifier} taxonomy + recorder that instantiates it; the two anchors mutually compose per §11.4.209 reviewer MINOR-1). Extraction sources: module 27 lines 1212, 1343, 1396, 1482-1488, 1498, 1613, 1631-1634; module 34 lines 5695, 5711-5717, 5722-5730, 3117-3168, 3180, 3201, 2313, 2381.

Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-240-PROPAGATION (literal 11.4.240) + recommended gates CM-PRODUCER-NOT-VERIFIER-SCOPE-SPLIT (a change's four scopes per (A) — producer, oracle, gate, verifier — resolve to distinct actors; a producer-authored verdict, a producer-modifiable gate config, a producer-modifiable oracle, or a producer-modifiable test

file in the same change violates the split) + CM-LEAST-PRIVILEGE-CAPABILITY-NOT-INSTRUCTION (the enforcement channel is a filesystem-scope / branch-protection / verdict-store-isolation / execution-isolation mechanism per (B), NOT an instruction / prompt / policy-document / social-contract; a project whose only enforcement is a written policy that "producers do not touch tests" fails the capability check) + paired §11.1 mutations (make the producer branch a modification to a tests/* file within the same change without a non-producer approver → the CODEOWNERS + capability check FAILs; author a status write of "verified: true" from the same actor that produced the diff, with no harness-written §11.4.115(F) verdict pair → the §11.4.146(D3) custody chain FAILs; grant the producer write access to the gate configuration + expect the review to catch the collapse → the capability-not-instruction gate FAILs on the enforcement-channel check; strip the literal → propagation FAILs; gate-code = separate work item, NOT claimed shipped §11.4.6/§11.4.227).

Canonical authority: constitution submodule Constitution.md §11.4.240. Non-compliance is a release blocker. No escape hatch — no --producer-may-write-tests, --self-certification-OK, --exhortation-suffices-as-enforcement, --producer-authored-verdict-accepted, --capability-limit-optional, --policy-document-is-enforcement, -> single-actor-for-all-scopes flag.

§_{26 08 15} . . — Illegal-state-unrepresentability preference: types before API-shape before lint before property-test before runtime-assertion (research-derived, — —)

Full anchor text per Phase 3 landing brief. See

curriculum_anchors_241_254_full_drafts.md §11.4.241 canonical block (Part A). Anchor establishes the closed ordered ladder (1) TYPE SYSTEM → (2) API SHAPE → (3) LINT → (4) PROPERTY-BASED TEST → (5) RUNTIME ASSERTION for every load-bearing invariant; a defect preventable at rung N caught only at rung > N is a finding. Composes §11.4.108 (pushes enforcement LEFT of Layer 1) / §11.4.115 / §11.4.142 / §11.4.194 / §11.4.201 / §11.4.224 / §11.1 / §11.4.6 / §11.4.35. Includes rung-by-language capability baseline (strongly-typed:

rungs 1-5; dynamically-typed with gradual: rung 1 conditional; untyped scripting: rungs 3-5 unconditional); genuinely-limited-environment clarification per §11.4.209 reviewer MINOR-4 (POSIX sh with only rung 5 SATISFIES the anchor IFF the runtime assertion GENUINELY CATCHES the invariant AND is paired with its §1.1 meta-test mutation per §11.4.107(10) self-validation). Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-241-PROPAGATION` (literal `11.4.241`) + recommended gate `CM-ILLEGAL-STATE-LADDER-PREFERENCE` + paired §1.1 mutation (gate-code = separate work item, NOT claimed shipped §11.4.6/§11.4.227).

Canonical authority: constitution submodule `Constitution.md` §11.4.241. Non-compliance is a release blocker. No escape hatch — no `--runtime-assert-suffices`, `--skip-type-preference`, `--dynamic-language-exempts-all-rungs`, `--property-test-in-lieu-of-type`, `-1→decorative-type-annotation-OK` flag.

2026 08 15

§ . . — Bisection-not-blame regression cause-finding + failure modes (research-derived, - -)

Full anchor per Phase 3 landing brief §11.4.242. When a previously-working feature/behaviour regresses, the FIRST-choice localisation mechanism is BISECTION (`git bisect` or equivalent monotone-search, $O(\log N)$) over the ordered commit range between the last known-good tag per §11.4.114 and current HEAD, NOT blame-log inspection, NOT speculative root-cause hunting. MONOTONE-SEARCH INVARIANT: bisection valid only when observable is monotone across the range (broken stays broken); non-monotone (flaky, environment-dependent, timing-race) MUST stabilise first — bisection over flaky signal produces a WRONG "guilty" commit with the same confidence as a correct one (§11.4/§11.4.1 bluff at the localisation layer). FAILURE-MODE CATALOGUE (closed 6-class set): (1) flaky/non-monotone observable, (2) build-broken commits mistaken for "bad" (use `bisect skip` exit 125), (3) merge-commit introductions (defect in conflict resolution not either parent), (4) external/environmental changes, (5) delayed-fuse introductions (latent commit + later trigger), (6) test-infra regressions masquerading as product regressions. `git bisect run` script MUST use exit-code discipline (0=good, 1-124=bad, 125=skip, 128+=abort); mapping

every non-zero to "bad" collapses the taxonomy. Bisection RESULT is captured evidence per §11.4.5/§11.4.69 feeding §11.4.102 systematic-debugging's root-cause phase. Honest boundary §11.4.6 — bisection localises the INTRODUCING COMMIT of a MONOTONE regression, does NOT prove that commit's change is the ROOT CAUSE. Composes §11.4.5 / §11.4.6 / §11.4.69 / §11.4.102 / §11.4.114 / §11.4.115 / §11.4.201 / §11.4.224 / §1.1. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-242-PROPAGATION (literal 11.4.242) + recommended gate CM-BISECTION-BEFORE-BLAME + paired §1.1 mutation.

Canonical authority: constitution submodule Constitution.md §11.4.242. Non-compliance is a release blocker. No escape hatch — no --blame-log-suffices , --skip-bisection , --flaky-observable-bisect-anyway , --run-script-exit+code-collapse-OK , --regression-guess-without-localisation flag.

2026 08 15

§ . . — Characterization / golden-master safety net for legacy code before behaviour-changing modification (research-derived, — —)

Full anchor per Phase 3 landing brief §11.4.243. Before ANY behaviour-changing modification to legacy code lacking characterization tests (code inherited without written behavioural contract, code whose original author has moved on, code whose documented behaviour has diverged from actual behaviour, code that "works but nobody knows why"), the modifier MUST FIRST author CHARACTERIZATION TESTS asserting CURRENT observable behaviour as GROUND TRUTH (whether "correct" or not) and land them as a green baseline. Only AFTER the safety net exists may the modification proceed. Pattern originates in Feathers's WORKING EFFECTIVELY WITH LEGACY CODE + generalises GOLDEN-MASTER (record outputs across representative input space; assert future runs match). LEGACY SCOPE (§11.4.6-honest): (i) lacks tests asserting observable behaviour; OR (ii) written spec diverges from measured actual behaviour; OR (iii) predates the constitution's TDD discipline (§11.4.224); OR (iv) modifier cannot describe its exact behaviour from reading source in time budget. AUTHORING LOOP: (1) identify observable surface → (2) generate representative input space → (3) run + record actual outputs as ground-truth baseline → (4) assert recorded values in tests → (5) commit characterization tests as SEPARATE commit

BEFORE the behaviour-changing modification. **DISTINCTION FROM CORRECTNESS:** characterization tests capture WHAT IS not WHAT SHOULD BE — DESCRIPTIVE not prescriptive. If current legacy behaviour is subtly wrong, the characterization test CAPTURES that wrongness; the modifier's job is to notice, ASK the operator per §11.4.66 whether to preserve or correct that behaviour, never silently. Composes §11.4.115 (RED-first — characterization is legacy counterpart) / §11.4.124 (investigate-before-remove — parallel discipline) / §11.4.118 (sampled input space = honest gaps, not completeness) / §11.4.201 / §11.4.224 (test-first backfilled for legacy). Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-243-PROPAGATION` (literal `11.4.243`) + recommended gate `CM-CHARACTERIZATION-BEFORE-LEGACY-MODIFICATION` + paired §1.1 mutation.

Canonical authority: constitution submodule `Constitution.md` §11.4.243. Non-compliance is a release blocker. No escape hatch — no `--modify-untested-legacy-directly`, `--skip-characterization`, `--claim-tests-not-needed-because-obvious`, `--decorative-assertion-suffices`, `-i refactor-without-safety-net` flag.

2026 08 15

§ . . — Cross-boundary contract tests + can-i-deploy gate (research-derived, - -)

Full anchor per Phase 3 landing brief §11.4.244. Changes modifying a boundary between independently-deployable components (service ↔ consumer, library ↔ dependents, DB schema ↔ readers/writers, message topic ↔ subscribers) MUST be covered by CONSUMER-DRIVEN CONTRACT TESTS on both sides AND a `can-i-deploy` gate MUST verify proposed provider version compatibility against every currently-deployed consumer version BEFORE the change ships. §11.4.244 MERGES contract-tests + can-i-deploy (they are useless in isolation — contract tests without the gate produce compatibility knowledge nobody consults; a can-i-deploy gate without contract tests has nothing to check). CDC pattern: consumers publish assumptions they make about the provider (Pact / Spring Cloud Contract / gRPC schema-evolution rules); providers verify each consumer's contract against actual implementation; provider change breaking a live consumer's contract fails provider's build BEFORE the change ships — "a precise red at the source." CAN-I-DEPLOY: at every deploy attempt, gate queries the contract broker

for the compatibility matrix; single red cell BLOCKS the deploy with machine-readable verdict naming the incompatible pair, NEVER warning-with-continue. MISSING contract = first-class refusal reason per §11.4.201 (a false-positive REFUSAL is a FAIL-bluff, false-negative allowance is a PASS-bluff). SemVer becomes MEASURABLE (PATCH bump breaking contract = mis-labelled version). Backward-compatibility window declared per §11.4.35. Composes §11.4.5 / §11.4.6 / §11.4.35 / §11.4.66 / §11.4.108 (per-side four-layer) / §11.4.135 (contract tests as permanent regression guards) / §11.4.185 (composed-system manual QA still needed) / §11.4.201 (gate honesty) / §1.1. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-244-PROPAGATION (literal 11.4.244) + recommended gates CM-CONTRACT-TESTS-BOTH-SIDES + CM-CAN-I-DEPLOY-GATE-WIRED + paired §1.1 mutations.

Canonical authority: constitution submodule Constitution.md §11.4.244. Non-compliance is a release blocker. No escape hatch — no --skip-contract-tests, --deploy-without-compatibility-check, --assume-compatible-when-broker-silent, --patch-may-break-contract, --contract-tests-optional-for-internal-boundaries flag.

2026 08 15

§ . . — Oracle-problem-first test
authoring: identify the oracle BEFORE authoring
the test (research-derived,
— —)

Full anchor per Phase 3 landing brief §11.4.245. Every test MUST FIRST identify its ORACLE — the mechanism deciding PASS vs FAIL for the input it drives — and the oracle MUST be DIFFERENT from the code under test AND from the test's own assertion mechanism. Authoring without explicit oracle identification produces the "test agrees with the code" anti-pattern (test asserts what the code happens to produce, so it passes as long as the code produces something, never catches wrong-output — the most common form of §11.4/§11.4.1 bluff at the test-authoring layer). The oracle problem is the OLDEST unsolved problem in software testing; §11.4.245 does not solve it, it mandates HONESTLY IDENTIFYING which oracle strategy each test uses so the reviewer per §11.4.194 can judge whether the identified oracle is genuinely independent. ORACLE STRATEGY CLOSED SET (each test's oracle NAMED explicitly in source): (1) SPECIFIED (spec/RFC/

standard names expected output); (2) DERIVED (independent reference implementation); (3) METAMORPHIC (metamorphic relations between related inputs — §11.4.107(8) generalised); (4) GOLDEN-MASTER / CHARACTERIZATION (past captured behavior — §11.4.243); (5) INVARIANT (universal property must hold — sort output sorted + permutation of input); (6) STATISTICAL (distribution properties for stochastic outputs, never specific value); (7) HUMAN (§11.4.185 manual QA). NO other classes. STRUCTURAL INDEPENDENCE: oracle's SOURCE of expected value does not depend on code being tested. §11.4.107(10) SELF-VALIDATION generalised: paired §1.1 mutation must produce a WRONG output; oracle catches it. Composes §11.4.5 / §11.4.6 / §11.4.107(10) / §11.4.115 / §11.4.194 / §11.4.201 / §11.4.224 / §11.4.243 / §1.1. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-245-PROPAGATION (literal 11.4.245) + recommended gate CM-ORACLE-STRATEGY-NAMED-AND-INDEPENDENT + paired §1.1 mutation.

Canonical authority: constitution submodule Constitution.md §11.4.245. Non-compliance is a release blocker. No escape hatch — no --test-agrees-with-code, --skip-oracle-identification, --oracle-is-the-code, --assertion-is-oracle, --implicit-oracle-OK flag.

2

2026 08 15

§ . . – Reproducible + hermetic builds + supply-chain integrity at SLSA Build Level (operator decision, – –)

Full anchor per Phase 3 landing brief §11.4.246. Every consuming project's build pipeline MUST produce REPRODUCIBLE, HERMETIC artifacts and MUST meet SLSA Build Level 2 as the FLEET-WIDE MINIMUM (operator decision, 2026-08-15 — Level 3+ tracked as backlog per §11.4.197). REPRODUCIBLE = same source at same time on any conformant host produces byte-identical artifacts (or diff-set fully accounted for by declared non-determinisms pinned to canonical values). HERMETIC = no undeclared inputs (every dep version-pinned + lockfile-hashed, every tool version declared, no unpinned public-network reach at build, no host-env leakage). SLSA-L2 = signed provenance describing source URI + commit hash + builder identity + invocation + artifact hash, on hosted tamper-resistant build platform (boba already satisfies ancillary: SSH-only §2.1, rootless podman

§11.4.161, fixed-tag pins §11.4.30/§11.4.77). Every consumer MAINTAINS docs/security/SLSA_LEVEL.md (or per §11.4.35 equivalent) recording CURRENT L level with cited evidence + TARGET L + tracked §11.4.197 upgrade items — CURRENT below L2 without tracked item is violation. SUPPLY-CHAIN EXTENSION: deps either vendored (checked in hash-verified), from provenance-attested source (npm/PyPI/crates.io as they add SLSA), OR mirrored through hash-pinning local registry. Threat model NARROW: covers source → artifact chain, does NOT cover malicious commit landing in source (§11.4.142 review) nor compromised dep source nor bug-free artifact (§11.4.108 four-layer + §11.4.185 manual QA). Composes §2.1 / §11.4.6 / §11.4.30 / §11.4.35 / §11.4.108 / §11.4.113 / §11.4.135 / §11.4.142 / §11.4.161 / §11.4.184 / §11.4.185 / §11.4.197 / §11.4.209 / §1.1. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-246-PROPAGATION (literal 11.4.246) + recommended gates CM-BUILD-REPRODUCIBILITY + CM-BUILD-HERMETICITY + CM-SLSA-L2-PROVENANCE-VERIFIED + CM-SUPPLY-CHAIN-DEPS-PROVENANCE-ATTESTED + paired §1.1 mutations.

Canonical authority: constitution submodule Constitution.md §11.4.246. Non-compliance is a release blocker. No escape hatch — no

```
--skip-reproducibility , --allow-network-during-build , --no-provenance , --slsa-optional , --below-l2-OK ,
--unpinned-dependency-OK , --host-env-may-leak-into-build flag.
```

2026 08 15

§ . . – Composite-output layer-move must move every layer (research-derived, – –)

Full anchor per Phase 3 landing brief §11.4.247. Changes MOVING a system-boundary layer (service migrates hosts, endpoint URL changes, API version bumps, DB schema shifts, message topic renames, ingress path re-routes, storage bucket moves) MUST move EVERY LAYER of the composite output atomically or via explicit compatibility window — NEVER piecemeal. LAYER TAXONOMY per §11.4.35 baseline: (i) server (handler + route + OpenAPI/proto/gRPC descriptor); (ii) client (call-sites + SDK stubs + integration tests); (iii) infrastructure (LB + ingress + API gateway + CDN + service-mesh + DNS); (iv) data stores (cached keys + DB URL rows + MQ subscriptions); (v) observability (metric labels + span names + log

fields + dashboards + alerts + SLOs); (vi) documentation (user + dev + runbook + README + cookbook); (vii) tests (integration + e2e + contract per §11.4.244 + load + smoke). Consumer provides `scripts/audit/layer_move_completeness.sh` reporting residual references delta; change proposing move MUST show ZERO residuals OR explicit compatibility-window plan per residual with tracked §11.4.197 items. COMPATIBILITY-WINDOW: old value keeps working alongside new until stated date/version/event, deprecation warnings named per caller, tracked deadline, clean removal at expiry — no "clean up later." [MATERIAL-THIN: Module 32 supplies the composite-output pattern (multiple co-varying layers, atomic move discipline, surface-swap-not-recreate); the specific web/service layer inventory is generalised from the pattern's shape per §11.4.17 with consumer §11.4.35 DATA supplying the concrete inventory.] Composes §11.4.6 / §11.4.35 / §11.4.111 / §11.4.118 / §11.4.124 / §11.4.135 / §11.4.191 / §11.4.197 / §11.4.238 (residual = coverage escape) / §11.4.244 / §1.1. Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-247-PROPAGATION` (literal `11.4.247`) + recommended gate `CM-LAYER-MOVE-COMPLETENESS-CHECK` + paired §1.1 mutation.

Canonical authority: constitution submodule `Constitution.md` §11.4.247. Non-compliance is a release blocker. No escape hatch — no `--partial-layer-move-OK`, `--skip-completeness-check`, `--residuals-cleaned-up-later`, `--observability-layer-optional`, `--doc-move-can-lag` flag.

2026 08 15

§ . . — Flaky-test quarantine + protected-regression-spec gate via [PROTECTED-SPEC: ATM-NNN] tag + CODEOWNERS required reviewer (operator decision, — —)

Full anchor per Phase 3 landing brief §11.4.248. Flaky tests are the SILENT ANTI-BLUFF DECAY vector (every flake conditions team to ignore red until real regression is dismissed as "probably flaky, re-run"). Two disciplines bound: (A) QUARANTINE — mechanical detection (test-history aggregator flags two-consecutive-run-different-verdict OR rerun-after-red pattern) moves flaky test to `tests/quarantine/` (per §11.4.35), opens tracked §11.4.197 stabilisation deadline, blocking suite runs WITHOUT quarantined test so green means green; deadline past = §11.4.66 extension or violation; (B) PROTECTED-REGRESSION-

SPEC — regression tests codifying known-broken defect fixes (per §11.4.135 permanent-guard suite) MAY be tagged `# [PROTECTED-SPEC: ATM-NNN]` (target-language comment equivalent) referencing the workable-item ATM-NNN (§11.4.54) they guard, marking test as load-bearing regression guard whose modification requires elevated review. Tag applied deliberately per test (not blanket — that voids the discipline). CODEOWNERS-REQUIRED-REVIEWER GATE: consuming project's CODEOWNERS declares modifications to `[PROTECTED-SPEC]` -tagged tests OR `regression/**/*` path scope require designated required-reviewer approval BEFORE merge; composes with §2.1 SSH + §11.4.113 no-force-push (gate cannot be silently bypassed); audit trail in PR review thread + commit history. Gate CATCHES change so reviewer distinguishes STRENGTHEN (approve) vs FIX-DEFECT-IN-TEST (approve after verifying ATM-NNN's defect not silently re-introduced) vs WEAKEN/REMOVE (refuse or §11.4.66 escalate) — makes modification VISIBLE not preventable. Composes §2.1 / §11.4.6 / §11.4.35 / §11.4.50 (ideal empty quarantine) / §11.4.54 / §11.4.66 / §11.4.113 / §11.4.115 / §11.4.135 / §11.4.197 / §11.4.226 / §1.1. Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-248-PROPAGATION` (literal `11.4.248`) + recommended gates `CM-FLAKY-TEST-QUARANTINE-WIRED` + `CM-PROTECTED-REGRESSION-SPEC-GATE` + paired §1.1 mutations.

Canonical authority: constitution submodule `Constitution.md` §11.4.248. Non-compliance is a release blocker. No escape hatch — no `--rerun-until-green`, `--quarantine-permanent`, `--protected-spec-review-optional`, `--codeowners-skip`, `--blanket-protect-everything`, `--tag-without-atm-id` flag. ⁴ ₂₄₉

2026 08 15

§ . . — Producer ≠ oracle ≠ gate ≠ verifier + flight recorder (research-derived, — —)

Full anchor per Phase 3 landing brief §11.4.249. Every load-bearing quality mechanism has FOUR distinct roles that MUST be structurally separate — (1) PRODUCER (code computing/emitting/building output); (2) ORACLE (mechanism deciding correctness per §11.4.245); (3) GATE (seam consuming oracle verdict to allow/block/proceed); (4) VERIFIER (mechanism auditing gate's decision — paired §1.1 mutation proving gate can refuse a bad output). Collapsing any two produces

specific bluff class: producer=oracle (§11.4.245 oracle-agrees-with-code); oracle=gate (refuses to say "cannot decide" defaults to allow, §11.4.201(6) FALSE-NULL); gate=verifier (self-auditing gate — §11.4.142 self-review anti-pattern applied). FLIGHT RECORDER: mechanism CAPTURING producer inputs + oracle verdict + gate decision + verifier audit as durable replayable artifact (structured JSONL event stream per §11.4.116, content-addressed per §11.4.207) so post-hoc "why did this go through" has evidence not shrug. Recorder is NOT the gate — a gate reading its own recorder to decide collapses gate + verifier + recorder. DETECTABLE COLLAPSES reviewer per §11.4.194 scans for: (1) producer=oracle (test's assertion calls code being tested), (2) oracle=gate (no "cannot decide" branch, defaults PASS), (3) gate=verifier (meta-test IS the gate — mutation making gate always-pass ALSO makes meta-test always-pass; meta-test MUST live OUTSIDE gate), (4) recorder=decision-substrate (gate uses recorder contents as decision input — precedent-based gating not authorised). [MATERIAL-THIN: four-role taxonomy generalised from producer/verifier separation + oracle-problem discussion + gate mechanics; FLIGHT RECORDER derived from clean-checkout re-run + attached receipts pattern per §11.4.17.] Composes §11.4.5 / §11.4.6 / §11.4.107(10) / §11.4.116 / §11.4.142 / §11.4.194 / §11.4.201 / §11.4.207 / §11.4.240 (producer ↔ verifier separation principle — mutual composes per §11.4.209 reviewer MINOR-1) / §11.4.245 / §1.1. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-249-PROPAGATION (literal 11.4.249) + recommended gates CM-ROLE-SEPARATION-DECLARED + CM-FLIGHT-RECORDER-CAPTURED + paired §1.1 mutations.

Canonical authority: constitution submodule Constitution.md §11.4.249. Non-compliance is a release blocker. No escape hatch — no --collapse-producer-oracle , --gate-self-audits , --skip-flight-recorder , --oracle-defaults-pass , --decision-without-recorded-rationale flag.

§ . . – Heuristic-tower signals a primitive defect: when N heuristic layers stack to compensate, the primitive is broken (research-derived, – –)

Full anchor per Phase 3 landing brief §11.4.250. Codebase accumulating TOWER of heuristics compensating for same underlying issue (retry-with-backoff + connection-pool-warmup + DNS-cache-flush + health-check-with-startup-grace — all guarding one flaky DB connection) IS the diagnostic: PRIMITIVE is defective, every layer papers over a manifestation without fixing cause. Extraction line 6486: heuristic-compensating-for-broken-primitive taxonomy row + ATMOSphere

BAD_INDEX (6) compensator-tree case study. Recognition signals: (i) each layer justification "we added because sometimes X happens" — reactive; (ii) removing any layer INCREASES failures BUT keeping all shows occasional residual; (iii) no single layer owns responsibility; (iv) layers added at different times by different authors per specific incidents; (v) monotone layer-count growth over time — debugging session for residual adds layer N+1 rather than fixing 1..N. DISTINCT from defense-in-depth: defense-in-depth is deliberate + each layer targets DIFFERENT failure class + designed together not accreted + removing layer produces SPECIFIC different failure not "same failure more often." SYSTEMATIC-DEBUGGING RESPONSE on tower depth ≥ 2 : (1) STOP adding layers; (2) INVOKE §11.4.102 systematic-debugging on PRIMITIVE; (3) proceed root-cause / pattern / hypothesis / implementation phases; (4) fix primitive; (5) REMOVE layers one at a time verifying with §11.4.5 evidence; (6) commit each removal per §11.4.124 with git-history evidence citing primitive fix. Composes §11.4.5 / §11.4.6 / §11.4.102 / §11.4.124 / §11.4.194 / §11.4.201 (outer layer masks §11.4.201(6) false-null) / §11.4.226 / §1.1. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-250-PROPAGATION (literal 11.4.250) + recommended gate CM-HEURISTIC-TOWER-TRIGGERS-PRIMITIVE-DEBUG + paired §1.1 mutation.

Canonical authority: constitution submodule Constitution.md §11.4.250. Non-compliance is a release blocker. No escape hatch — no --add-another-heuristic-layer , --skip-primitive-debug , --tower-is-defense-in-depth-without-justification , --reactive-mitigation-OK , --layer-per-incident flag.

§ . . – Byte-identical-fork prohibition + role-as-data-pack (research-derived, – –)

Full anchor per Phase 3 landing brief §11.4.251. When project needs would be met by TWO NEAR-IDENTICAL CODE COPIES differing only in configuration/data/role-specific parameters (prod + staging with hardcoded hostnames; reader + writer with different flag values; web + mobile with different feature flags), copies MUST NOT be maintained as parallel forks. Instead ONE code copy retained + role-specific values extracted into DATA PACK (role-config file / env var set / capability descriptor / feature-flag bundle) consumed by the one code to instantiate as either role. Extraction line 6490: byte-identical-fork taxonomy row + SonarQube-Factory case (Application/main.kt byte-for-byte the caching-proxy launcher, still in package net.milosvasic.factory.proxy.caching.application — "no SonarQube code"; fix = role-real acceptance test, "roles become data-only packs you reference, never repos you fork"). DETECTION: reviewer or mechanical scan finds byte-identical fork when diff shows every difference is a config value (hostname/port/flag/label/env/credential-ref) NOT behavioral, common git-history ancestor, > 90% byte identity; per §11.4.35 threshold-DATA. EXTRACTION: (1) IDENTIFY configuration surface; (2) EXTRACT to data pack per role; (3) REFACTOR single code copy to consume at instantiation (constructor injection / factory / config-loader-at-startup); (4) DELETE one fork; (5) VERIFY per §11.4.5 both role instantiations produce same observable behavior their forks previously produced (§11.4.243 characterization before refactor); (6) commit extraction per §11.4.124 with git-history evidence. LEGITIMATE FORKS: temporary migration (tracked §11.4.197 end date), experimental / A-B (bounded end date), truly divergent behavior (copies differ in algorithm/protocol/data-model not just config). Composes §11.4.5 / §11.4.6 / §11.4.28 (decoupling — role-configurable code shared across consumers) / §11.4.35 (data-pack location) / §11.4.124 / §11.4.177 / §11.4.194 / §11.4.197 / §11.4.243 / §1.1. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-251-PROPAGATION (literal 11.4.251) + recommended gate CM-NO-BYTE-IDENTICAL-FORK + paired §1.1 mutation.

Canonical authority: constitution submodule Constitution.md §11.4.251. Non-compliance is a release blocker. No escape hatch — no --fork-for-role-OK , --skip-data-pack-extraction , --copy-and-modify-config , --duplicate-code-per-environment , --near-identical-forks-are-fine flag.

§ . . – Fail-closed-on-dangerous-combination for mutating / credential surfaces (research-derived, – –)

Full anchor per Phase 3 landing brief §11.4.252. Every code path COMBINING ≥ 2 dangerous capabilities MUST FAIL CLOSED — refuse operation unless every capability's precondition verifiably satisfied — NEVER fail open on ambiguous/missing/unresolvable input. Security-engineering generalisation of §11.4.201 guard-honesty applied to code paths where false-negative blast radius dwarfs operator inconvenience of false-positive. DANGEROUS-COMBINATION TAXONOMY: (1) MUTATION of shared resource (DB / filesystem / container state / config), (2) UNTRUSTED INPUT (user string / network bytes / external env var / unverified file), (3) CREDENTIAL ACCESS (secret read / authenticated API call / signing key use), (4) EXTERNAL SIDE EFFECT (email / webhook / MQ publish / payment endpoint), (5) SHELL / EXEC (subshell / binary with user-influenced args / template eval with user values), (6) IRREVERSIBLE (delete / drop / truncate — §11.4.113 force-push already absolutely forbidden). Any path combining ≥ 2 = dangerous-combination, MUST fail closed. INVARIANT: (1) VERIFY every precondition (authorization + input shape + credential presence + target identity); (2) REFUSE when ANY unverifiable with ERROR EXIT naming specific unresolved precondition; (3) EMIT captured evidence per §11.4.5 of refusal; (4) NEVER default to allow/proceed/retry-with-relaxed on ambiguity — §11.4.101 conservative-safe default = REFUSE. FAIL-OPEN ANTI-PATTERNS refused: `catch { /* ignore */ }, credential = credential || default, if (!valid) { log("warn"); return proceed(input); }, target = user_input || default_target, unbounded retry { dangerous(); }`. Distinct from graceful degradation. [MATERIAL-THIN: 6-class dangerous-combination taxonomy is a §11.4.17 generalisation of the fail-closed pattern's shape applied to specific dangerous surfaces; closed-set formalisation, not literal per-class enumeration.] Composes §11.4.5 / §11.4.6 / §11.4.10 (credentials specialisation) / §11.4.66 (operator escalation when refuse stuck) / §11.4.101 / §11.4.113 (max instantiation) / §11.4.194 / §11.4.201 / §1.1. Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-252-PROPAGATION` (literal `11.4.252`) + recommended gate `CM-DANGEROUS-COMBINATION-FAIL-CLOSED` + paired §1.1 mutation.

Canonical authority: constitution submodule `Constitution.md` §11.4.252. Non-compliance is a release blocker. No escape hatch — no `--fail-open-on-error` , `--swallow-exception` , `--default-missing-credential` , `--proceed-on-ambiguous-input` , `--retry-until-succeeds` flag.

2026 08 15

§ . . — Idempotency under retry + DB-level durable uniqueness guard (research-derived, — —)

Full anchor per Phase 3 landing brief §11.4.253. Every retryable operation (network calls with auto-retry, at-least-once MQ consumers, cron with overlapping schedules, HTTP endpoints with client-retry, DB transactions with retry-on-serialisation-failure, deploy pipelines with rerun capability) MUST BE IDEMPOTENT: N executions produce same observable end state as one. Idempotency ENFORCED at LOWEST LAYER where "same end state" is durably provable — typically DATABASE UNIQUENESS CONSTRAINT on caller-supplied idempotency key, NEVER only application-layer check-then-insert (races). Prevents DOUBLE-CHARGE / DOUBLE-EMAIL / DOUBLE-DELETE / DOUBLE-EFFECT bugs. IDEMPOTENCY-KEY DISCIPLINE: every retryable op accepts (or generates + passes forward) unique identifier naming THIS logical op; persistence layer has UNIQUE INDEX / CONSTRAINT on key column; second attempt REFUSED by DB at durable level not application logic reading-then-writing; refusal caught by caller as "duplicate — already succeeded" = SUCCESS. ANTI-PATTERNS: (i) check-then-insert without unique index races; (ii) app-level dedup table read without spanning-transaction; (iii) key scoped to app instance (in-memory not shared); (iv) effect persisted BEFORE key (crash between = double-effect on retry); (v) no key at all + "the effect is 'set X to V' so redoing is fine" (safe for pure sets, NOT for INSERTs / MQ emissions / external-system side effects); (vi) idempotency-key expires — converts to time-bounded-idempotent = §11.4.6 mis-labelling if caller expects unbounded. CHAOS VERIFICATION per §11.4.85: (1) SIMULTANEOUS RETRIES (same key from 2 concurrent callers), (2) MID-OPERATION CRASHES (kill receiver between "check exists" and "insert effect"), (3) RETRY AFTER APPARENT SUCCESS (caller believes fail, receiver persisted, retry MUST recognise duplicate). External-system side effects (Stripe/PayPal charging card + local DB write): idempotency key MUST pass through to external; if external does not support, op is not fully

idempotent — label honestly per §11.4.6. Composes §11.4.5 / §11.4.6 / §11.4.85 / §11.4.201 / §11.4.252 / §1.1. Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-253-PROPAGATION` (literal `11.4.253`) + recommended gates `CM-IDEMPOTENCY-DB-LEVEL-UNIQUE-GUARD` + `CM-IDEMPOTENCY-CHAOS-VERIFIED` + paired §1.1 mutation.

Canonical authority: constitution submodule `Constitution.md` §11.4.253. Non-compliance is a release blocker. No escape hatch — no

`--check-then-insert-OK`, `--application-dedup-suffices`, `--skip-idempotency-chaos-test`, `--claim-idempotent-without-key`, `--expire-idempotency-marker-silently` flag.

2026 08 15

§ . . — Boot-time invariant assertion + capability matrix (research-derived, - -)

Full anchor per Phase 3 landing brief §11.4.254. Every service / process / long-running executable MUST assert BOOT-TIME INVARIANTS at startup — required env vars present + valid; required config files parseable + declared schema version matches; required deps reachable; required credentials present + not-expired; required schema migrations applied; required feature-flag states declared; required host resources present (disk / memory / thread limits per §12.12 / port availability) — and FAIL FAST + LOUD if any violated, BEFORE accepting request / performing work. Extraction line 6488 (Catalogizer SQLCipher case: `catalog.db` 's first 16 bytes are ASCII SQLite format 3 — plaintext despite "encrypted at rest"; fix = boot-time-check-that-fails-closed verifying `cipher_version` at startup) + line 6489 (stubbed-core-behind-claim + explicit capability matrix with every adapter reports `IsImplemented()` /attempted-vs-not + challenge that fails against no-op stub). Prevents "service silently starts with broken invariant + serves requests that fail deep in request path hours after boot." INVARIANT CLASSES: (i) environment, (ii) configuration, (iii) dependencies, (iv) credentials, (v) schema, (vi) capabilities, (vii) host. FAIL-FAST-LOUD: (1) EMIT structured error to stderr + captured log naming specific failed invariant; (2) EXIT non-zero (traditionally 78 `EX_CONFIG` for config, consumer supplies per §11.4.35); (3) DO NOT accept requests / begin main event loop / bind listening port. Service booting partially + serving requests-that-will-fail is worse than refusing to start — partial-boot masquerades healthy long enough for

load balancer to route traffic. Dual to §11.4.252 fail-closed on dangerous combinations. CAPABILITY MATRIX: consumer's docs/services/CAPABILITY_MATRIX.md (per §11.4.35) declares per service (i) boot-time invariants, (ii) runtime capabilities OFFERED, (iii) runtime capabilities REQUIRED from other services. Regenerated / verified on every commit per §11.4.106 docs-chain + §11.4.86 fingerprint. Undeclared endpoints emitted = violation; silent upstream deps not declared = violation (invisible to §11.4.244 contract testing). Composes §11.4.5 / §11.4.6 / §11.4.35 / §11.4.86 / §11.4.106 / §11.4.108 Layer 3 (boot IS runtime-on-clean-target check) / §11.4.128 / §11.4.244 (contract tests consume matrix) / §11.4.252 / §12.12 / §1.1. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-254-PROPAGATION (literal 11.4.254) + recommended gates CM-BOOT-TIME-INVARIANT-ASSERTED + CM-CAPABILITY-MATRIX-PRESENT-AND-IN-SYNC + paired §1.1 mutations.

Canonical authority: constitution submodule Constitution.md §11.4.254. Non-compliance is a release blocker. No escape hatch — no --skip-boot-time-assertion , --boot-partial-OK , --default-missing-env-silently , --capability-matrix-optional , --service-with-undeclared-dependency-OK flag. ^{4 255}

2026 06 25

§ . . — Mandatory HelixTranslate canonical translation pipeline (User mandate, — —)

Forensic anchor — verbatim user mandate (2026-06-25): "Every project that inherits this constitution MUST use HelixTranslate ([git@github.com:HelixDevelopment/HelixTranslate.git](https://github.com/HelixDevelopment/HelixTranslate)) — via a universal, reusable pipeline — to generate ANY translations it needs. This is a hard rule/constraint."

ALL translations of ANY content — UI strings, documents, articles, CVs, READMEs, marketing copy, in-app messages, release notes, legal text, ANY human-language string rendered in a non-source language — MUST be produced by **HelixTranslate** ([git@github.com:HelixDevelopment/HelixTranslate.git](https://github.com/HelixDevelopment/HelixTranslate)) driven through the **canonical universal translation pipeline** shipped in this submodule at scripts/translation/ (translate-pipeline.sh + render-articles.sh). This is a hard, non-negotiable constraint.

The following are FORBIDDEN and constitute a §11.4.255 violation:

(a) **Hand-translation** — no human-authored, manually edited, or "good enough" hand-written translation of any content, ever. The engine is the single source of every translated string.

(b) **Other engines** — no Google Translate, DeepL, raw ChatGPT/Claude prompting, OS translation APIs, library i18n auto-translators, or any engine other than HelixTranslate's `unified-translator`. There is exactly one engine.

(c) **Silent fallback** — no quiet substitution of the source string, no empty/partial output passed off as a translation, no swallowed engine error, no "absence-of-error" PASS (§11.4 / §11.4.1). On exhaustion of all configured providers the pipeline MUST FAIL LOUD (non-zero exit, destination NOT written) so the caller detects it; a missing translation is a hard error, never a degraded-but-shipped string. Provider-to-provider failover WITHIN HelixTranslate (primary → fallback, with logged, evidenced retries) is NOT a silent fallback and is permitted; falling back to ANY non-HelixTranslate source is forbidden.

Mandatory validation. Every generated translation MUST be validated for accuracy with real (physical) captured evidence per §11.4.5 / §11.4.69 / §11.4.107 — no metadata-only / config-only / grep-without-runtime PASS (§11.4 / §11.4.1), no false results, no bluff of any kind (§11.4.6). The per-file run log (provider used, byte counts, retries, success/failure) is the evidence artifact; an unvalidated or unevidenced translation is treated as ABSENT.

Explicit provider routing. Provider + model routing MUST be explicit and recorded — never implicit/defaulted-in-the-dark. The canonical routing is provider `groq` / model `llama-3.3-70b-versatile` with documented failover to provider `mistral` / model `mistral-large-latest`; the provider that actually produced each output MUST be recorded in the per-file log. Any deviation from the canonical routing MUST be stated explicitly in configuration and logged.

Canonical pipeline location. `scripts/translation/translate-pipeline.sh` (universal Markdown / plain-text + article-frontmatter-preserving translator) and `scripts/translation/render-articles.sh` (article-source → HTML fragment renderer). Both are reusable verbatim across projects; the engine binary is supplied by the consuming project via `HELIX_TRANSLATE_BIN` and built from the HelixTranslate repository.

Classification: universal (§11.4.17) — a platform-neutral, language-neutral translation discipline reusable by ANY project; the consuming project supplies its concrete source files, target languages, and the HelixTranslate engine build per §11.4.35. Composes §11.4 / §11.4.1 / §11.4.5 / §11.4.6 / §11.4.7 / §11.4.69 / §11.4.107. Propagation gate `CM-COVENANT-114-255-PROPAGATION` (literal `11.4.255`) + recommended gate `CM-MANDATORY-HELIXTRANSLATE-PIPELINE` + paired §1.1 meta-test mutation (gate-code = separate work item).

Rationale. A single, evidenced, fail-loud translation engine guarantees that every translated string in the fleet has a known provenance, a reproducible pipeline, and a captured accuracy proof — eliminating the silent drift, inconsistent quality, and untraceable "where did this string come from" failures that hand-translation and ad-hoc engines produce. One engine, one pipeline, one evidence trail.

Canonical authority: this Constitution.md §11.4.255 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--allow-hand-translation`, `--other-engine-OK`, `--skip-translation-validation`, `--silent-fallback` flag exists.

2026 06 25

§ . . — Mandatory independent per-language translation review (User mandate, - -)

Forensic anchor — verbatim user mandate (2026-06-25): "Quality of translations for all languages MUST BE ALWAYS properly checked and reviewed! We MUST CREATE independent, specialized translations review agents for all languages we support which will be performing full testing, review, validation and verification for all translated content! This MUST BE mandatory part of our root constitution (Submodule) and followed as mandatory constraint / rule by this and all our projects implementing the constitution Submodule."

Every translated artifact produced under §11.4.255 MUST, before it is shipped / committed / deployed, be INDEPENDENTLY REVIEWED for quality by a per-language review agent. This is a hard, non-negotiable gate that composes with — and never replaces — the §11.4.255 generation mandate.

Requirements:

(a) **Independence.** The reviewer MUST be a different model/provider than the one that produced the translation (e.g. translator `mistral/mistral-large-latest` → reviewer `groq/llama-3.3-70b-versatile` , or any distinct verified model). A model reviewing its own output is NOT an independent review and is forbidden.

(b) **Per-language specialization.** The reviewer MUST act as a NATIVE speaker / linguistic reviewer of the SPECIFIC target language, evaluating that language's grammar, orthography, and correct script (Cyrillic for ru/be/kk, Arabic for ar/fa, Han/Kana/Hangul/Devanagari for zh/ja/ko/hi, Latin for sr/de/es/fr/tr, etc.).

(c) **Full criteria.** Each review MUST judge: accuracy/fidelity, fluency/naturalness, completeness (nothing omitted or invented), correct script & orthography, and absence of untranslated source leftovers (proper nouns / URLs / code excepted).

(d) **Strict machine-readable verdict + evidence.** Each review MUST emit a structured PASS/FAIL verdict with per-criterion scores and an issues list, persisted as a physical evidence artifact (§11.4.5 / §11.4.69 / §11.4.107). A FAIL — or an ERROR / unavailable-reviewer — BLOCKS the content; only a real PASS may ship. No metadata-only / absence-of-error PASS (§11.4 / §11.4.1); no bluff (§11.4.6).

(e) **Remediation loop.** A FAIL routes the artifact back through the §11.4.255 pipeline (re-translate / multipass-polish) and is re-reviewed, iterating until PASS — never hand-patched to pass (hand-translation remains forbidden under §11.4.255).

Canonical tooling. `scripts/translation/review-translations.sh` (per-language review driver) + `scripts/translation/review_translation.py` (independent reviewer engine emitting strict JSON verdicts). Reusable verbatim across projects; the consuming project supplies its source/translation file pairs and a reviewer provider DISTINCT from its translator.

Classification: universal (§11.4.17) — a platform-neutral, language-neutral translation-quality-assurance discipline reusable by ANY project that ships translated content; composes §11.4.255 (the translation it gates) / §11.4 / §11.4.1 / §11.4.5 / §11.4.6 / §11.4.69 / §11.4.107 / §11.4.134 (rock-solid evidence).

Propagation gate `CM-COVENANT-114-256-PROPAGATION` (literal `11.4.256`) + recommended gate `CM-MANDATORY-TRANSLATION-REVIEW` + paired §1.1 meta-test mutation (gate-code = separate work item).

Rationale. §11.4.255 guarantees a translation's provenance and a fail-loud pipeline, but not its linguistic quality. An independent, native-level, per-language reviewer with a captured PASS/FAIL verdict is what turns "a string was produced" into "a string was verified correct" — making translation quality an auditable fact rather than an assumption.

Canonical authority: this Constitution.md §11.4.256 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no

`--skip-translation-review` , `--self-review-OK` , `--accept-unreviewed-translation` flag exists.

§ . . — Comprehensive user-manual + guide + FAQ coverage: every component / service / ²⁰²⁶feature / user-visible workflow ships an operator-usable manual, task-oriented guide, and FAQ, always in sync, no orphan capability (BACKGROUND :: REMINDER :: IMPORTANT operator mandate, — —)

Verbatim operator mandate (2026-08-15, Point 1): *"Whole project and all work done MUST BE fully covered with full user manuals and guides and properly created FAQs."*

Compact summary: every consuming project MUST cover the WHOLE project (every component, service, subsystem, binary, container, plugin, feature, and user-visible workflow) with three cooperating documentation classes — (a) a complete USER MANUAL per component/service/product, (b) task-oriented GUIDES for every user-facing workflow, and (c) properly-created FAQs derived from real operator/QA/end-user questions — always in sync with the shipped code, always exported per §11.4.65, always reachable from README per §11.4.212, always regenerated via Docs-Chain per §11.4.106, and always tracked as workable items when incomplete per §11.4.197; a component / service / feature that ships without its trio of docs (manual + guide + FAQ) is a §11.4.257 violation of §11.4 PASS-bluff severity at the documentation layer (the feature ships to a user who cannot use it because nothing tells them how; the §11.4 anti-bluff covenant extended to

documentation completeness). §11.4.257 GENERALISES §11.4.18 (script docs — a single class of executable artifact) and §11.4.153 (per-feature Status set — a status not a manual) into the full three-class documentation floor for every user-facing surface, and BINDS §11.4.212 (README entry point) + §11.4.106 (Docs-Chain sync) + §11.4.65 (four-format export) into the always-in-sync guarantee.

(A) COMPREHENSIVE COVERAGE — no orphan capability. Every component / service / subsystem / binary / plugin / user-visible feature enumerated in the project's shipped surface (the §11.4.153 feature Status roster or its equivalent SSoT) MUST map to at least ONE user manual (component-level operator reference), at least ONE task-oriented guide per user workflow that surface exposes, and one FAQ entry per real recurring question captured from operators / QA / end-users. A capability the project ships that no manual documents, no guide covers, and no FAQ addresses is UNSHIPABLE regardless of how well the code works — the user has no path to use it.

(B) THREE DISTINCT DOCUMENT CLASSES, none substitutes another. (a) USER MANUAL — the component's operator reference: what it is, what it does, its inputs/outputs/env vars/config/ports/dependencies, its start/stop/status commands, its logs/metrics/health surface, its failure modes and their remediation; complete enough that an operator can run + supervise + debug it without reading the source. (b) GUIDES — task-oriented, workflow-shaped documents, one per user-facing workflow (setup, configuration, integration, migration, troubleshooting, security-hardening); each guide walks the user through a real journey, with prerequisites, step-by-step commands, expected output at each step, verification checkpoints, and rollback path. (c) FAQs — questions PROPERLY CREATED from real recurring signal (operator questions, QA reports, end-user tickets, agent-found-defect narratives), NEVER fabricated marketing FAQ; every entry cites the question source class + resolution + link to the manual/guide section that owns the durable answer.

(C) ALWAYS IN SYNC. Every doc carries the §11.4.44 revision header; Docs-Chain (§11.4.106) is bound at commit / build / constitution-pull write-seams so an edited component whose docs are stale REFUSES the commit; four-format export per §11.4.65 (.md + .html + .pdf, + .docx where the class warrants — the §11.4.153 exception); every doc reachable transitively from README per §11.4.212. FAQ

entries that describe now-fixed defects are RETAINED with a "resolved in vN.N.N" note (they answer the user asking "why did this happen to me"), never silently deleted.

(D) TRACKABLE COMPLETION. Missing docs are §11.4.197 items (research/kicked-off-work completion mandate) — an undocumented capability is an un-wired requirement whose loss is FORBIDDEN; the workable-items DB (§11.4.93/.95) carries per-surface docs-status rows so "what is documented / what is not" is a live query, never a periodic manual audit.

(E) HONESTY. Where a genuine gap exists (a new capability whose docs are in flight), it is HONESTLY MARKED at both the README entry point + the component's Status.md (§11.4.153) as `[DOC-PENDING: <tracker-id>]` per the §11.4.223 provenance-marker discipline, NEVER silently omitted, NEVER faked with a placeholder paragraph that says nothing (a fake-doc is a §11.4 PASS-bluff at the documentation layer — an operator reads it, learns nothing, and thinks they read the doc).

Honest boundary (§11.4.6). §11.4.257 mandates the DOCUMENTATION EXISTS, is IN SYNC, and is REACHABLE; it does NOT verify the doc is CORRECT in every semantic detail — that stays §11.4.99 (latest-source cross-reference) + §11.4.194 (exhaustive review including doc content) + §11.4.238 (any operator-found doc gap is a coverage-escape triggering a §11.4.138-family audit). "Comprehensive" is bounded to the shipped-surface enumeration; an internal implementation detail with no user-facing exposure has no user-manual obligation (documentation of it lives under §11.4.18 script/source class or design docs, not this anchor).

Classification: universal (§11.4.17). Composes §11.4.11 / §11.4.18 / §11.4.44 / §11.4.65 / §11.4.86 / §11.4.99 / §11.4.106 / §11.4.148 / §11.4.153 / §11.4.194 / §11.4.197 / §11.4.208 / §11.4.212 / §11.4.223 / §11.4.238. Propagation gate `CM-COVENANT-114-257-PROPAGATION` (literal `11.4.257`) + recommended gates `CM-USER-MANUAL-PER-COMPONENT` + `CM-GUIDE-PER-USER-WORKFLOW` + `CM-FAQ-FROM-REAL-QUESTIONS` + paired §1.1 meta-test mutations.

Canonical authority: constitution submodule `Constitution.md` §11.4.257. Non-compliance is a release blocker. No escape hatch — no `--skip-user-manual`, `--guide-later`, `--faq-optional`, `--capability-without-docs-OK`, `--fake-doc-placeholder-OK`, `--source-of-truth-is-code-comments` flag.

§ . . – Architectural + data-flow + state-machine + sequence-diagram coverage: every project ships accurate, exported, machine-derivable-or-honestly-authored diagrams embedded into the docs that describe the system (BACKGROUND :: REMINDER :: IMPORTANT operator mandate, – –)

Verbatim operator mandate (2026-08-15, Point 2): *"Whole project MUST BE fully covered with properly created graphs, diagrams, schemes which will be incorporated into all existing documentation."*

Compact summary: every consuming project MUST cover the WHOLE project with properly created graphs / diagrams / schemes across the FOUR canonical diagram classes — (a) SYSTEM ARCHITECTURE (services, containers, ports, dependencies, deployment topology), (b) DATA FLOW (how information moves between components, storage, external systems), (c) STATE MACHINE (lifecycle of every stateful entity: workable items, orchestrator sessions, long-ops, agents, deploys), (d) SEQUENCE (per critical user workflow and per critical cross-component protocol) — INCORPORATED into the existing documentation surface (README, user manuals per §11.4.257, guides, Status docs per §11.4.153, design docs per §11.4.218/§11.4.219) at the natural point of use, NEVER a lone diagram directory nobody opens; every diagram is (i) exported in a scale-appropriate open format (SVG for structural diagrams; PNG for raster fallback; the source in a diffable text form — Mermaid, PlantUML, D2, or a diffable SVG source — §11.4.220 open-first, never a proprietary-tool-only master), (ii) rendered non-blank + non-degenerate + accurate (the §11.4.107(10) analyzer discipline applied to diagrams — a rendered blank / mislabelled / truncated diagram is a §11.4/§11.4.1 bluff at the visual-doc layer), (iii) reachable from README per §11.4.212, and (iv) in sync via §11.4.106 Docs-Chain so a code change that invalidates a diagram REFUSES the commit until the diagram is updated. §11.4.258 STRENGTHENS §11.4.218 (design-library catalogue — components not architecture) and §11.4.219 (screen catalogues — UI IA not system topology) from design-side coverage into whole-project documentation-side diagram coverage.

(A) FOUR-CLASS COVERAGE. (a) ARCHITECTURE — the deployed / deployable topology: every service, container, binary, port, external dependency, host boundary, network boundary, storage boundary; kept true to the actual docker -

`compose.yml` / deployment manifest / running system. (b) DATA FLOW — how information (user request, event, credential, artifact, evidence) traverses the system: source → transformation → sink; each edge annotated with format + protocol + trust boundary. (c) STATE MACHINE — the lifecycle of every stateful entity the system manages: workable items (§11.4.15), long-ops (§11.4.232), orchestrator sessions (§11.4.187), agents (§11.4.147), deploy artifacts (§11.4.108), builds; transitions labelled with the triggering event and the guard that gates it. (d) SEQUENCE — one diagram per critical user workflow and per critical cross-component protocol; actors + messages + return paths + failure branches, not only the happy path.

(B) INCORPORATION — never orphan. Each diagram lives at the point of use (architecture diagram in README + `docs/ARCHITECTURE.md` ; per-service data-flow embedded in that service's user manual per §11.4.257; state-machine embedded in the doc that owns the state; sequence diagram in the guide for the workflow it depicts); a `docs/diagrams/` gallery may exist as an index, but the load-bearing appearance is EMBEDDED into the doc where the reader needs it, per §11.4.212 reachability.

(C) OPEN + DIFFABLE. Sources in diffable text (Mermaid / PlantUML / D2 / diffable SVG), NEVER proprietary-tool-only masters (§11.4.220); rendered outputs committed alongside sources so a reader without the render toolchain still sees the diagram; renderers §11.4.201-probed before use, absent-tool ⇒ §11.4.3 SKIP-with-reason not a fake output (§11.4.222 export-wave discipline applied here).

(D) ACCURACY VERIFICATION. Every diagram is rendered non-blank + non-degenerate + accurate: a probe render exists (a Mermaid parse-check; an SVG geometry check; a §11.4.107(10) golden-good vs golden-bad-with-drift fixture for the derived-from-code diagrams); a diagram whose depicted service list, port map, or state set diverges from the SSoT (compose file, workable-items schema, orchestrator registry) is caught by a §11.4.86-fingerprint diff on the SSoT + a stale-diagram audit run at pre-build; the Docs-Chain (§11.4.106) refuses a commit that changes an SSoT without an accompanying diagram refresh where a mapped diagram exists.

(E) DERIVED-WHERE-POSSIBLE. Prefer derived / generated diagrams over hand-drawn where the source data exists: a compose-file-driven architecture diagram, a schema-driven state machine, a route-table-driven sequence; hand-

drawn diagrams (for user-journey narratives, conceptual overviews) are LABELED as such with an authored-date and reconciled against the SSoT at each release round.

Honest boundary (§11.4.6). §11.4.258 mandates the diagrams EXIST, are ACCURATE against their SSoT, are EMBEDDED, and are RENDERED HONESTLY; it does NOT claim every possible diagram class is required for every project (an internal library with no runtime deployment has no ARCHITECTURE diagram in the deployment sense — it ships its module-dependency diagram in that class's place). Bounded to the WHOLE-PROJECT coverage the mandate names.

Classification: universal (§11.4.17). Composes §11.4.6 / §11.4.65 / §11.4.86 / §11.4.106 / §11.4.107(10) / §11.4.153 / §11.4.201 / §11.4.212 / §11.4.216 / §11.4.218 / §11.4.219 / §11.4.220 / §11.4.222 / §11.4.223 / §11.4.257. Propagation gate `CM-COVENANT-114-258-PROPAGATION` (literal `11.4.258`) + recommended gates `CM-ARCHITECTURE-DIAGRAM-PRESENT-ACCURATE` + `CM-DATAFLOW-DIAGRAM-PER-SERVICE` + `CM-STATEMACHINE-DIAGRAM-PER-STATEFUL-ENTITY` + `CM-SEQUENCE-DIAGRAM-PER-CRITICAL-WORKFLOW` + `CM-DIAGRAM-EMBEDDED-NOT-ORPHAN` + paired §1.1 mutations.

Canonical authority: constitution submodule `Constitution.md` §11.4.258. Non-compliance is a release blocker. No escape hatch — no `--skip-architecture-diagram`, `--diagram-in-orphan-directory`, `--proprietary-master-only`, `--stale-diagram-OK`, `--dataflow-implied-from-code`, `--sequence-later` flag.

§ . . – README quality-status badge row: at the top of every project README, a comprehensive row of quality badges with a closed green→amber→red vocabulary, each backed by machine-derived source data, illustrating quality and production-readiness in one glance (BACKGROUND :: REMINDER :: IMPORTANT operator mandate, – –)

Verbatim operator mandate (2026-08-15, Point 5): *"We MUST ADD at the top of the main README all possible badges with coloring from green (ok) to red (not ok) for all possible areas of work done to the project illustrating quality of work done and how close are we to deploying this project to full production."*

Compact summary: every consuming project MUST render, at the TOP of the main README (immediately below the H1, above the introduction — the reader's first visual signal), a comprehensive BADGE ROW carrying every applicable quality signal with a CLOSED color vocabulary from GREEN (ok) through AMBER (attention) to RED (not ok) — plus a "production-readiness" gauge that summarises how close the project is to full production deployment — with EVERY badge backed by MACHINE-DERIVED source data (never a hand-typed color-of-the-day, §11.4.6 no-guessing at the status-report layer); a project whose README lacks the badge row, or carries badges whose colors are fabricated / stale / not backed by real source data, is a §11.4.259 violation of §11.4 PASS-bluff severity at the reader-first-impression layer (an operator glances at the README and gets a green picture of a broken project). §11.4.259 EXTENDS §11.4.57 (README doc-link section) with a symmetric BADGE section at the same first-visible-position privilege, and BINDS §11.4.106 (Docs-Chain sync — badges regenerate on every state change) + §11.4.86 (fingerprint-driven sync — badge data derived from the SSoT the badge refers to) + §11.4.201 (guards must assert the real condition — a badge is a guard for the reader's eye).

(A) CLOSED COLOR VOCABULARY. GREEN = healthy at target (all constraints met, no findings); AMBER = attention (metric on-target but degrading, or off-target within acceptable band, or blocked on a tracked follow-up); RED = not ok (metric off-target beyond acceptable band, a §11.4.108 gate failing, a §11.4.135 uncovered guard, an open release-blocker); optional GRAY = N/A (badge class does not apply

to this project) with an inline note explaining why. NO other colors, NO other verdicts, NO "yellow" or "orange" or "beta" gradations — the closed vocabulary is the operator's cognitive contract with the badge row.

(B) BADGE CLASS FLOOR — every applicable area covered. The MINIMUM set every consuming project ships (adding project-specific badges as consumer-supplied §11.4.35 DATA is welcome, subtraction is not): (1) BUILD status (last build green/red, artifact fingerprint); (2) TEST-BREADTH per §11.4.27/§11.4.169 (unit/integration/e2e/full-automation/security/performance/anti-bluff — one badge per type, or a rollup with a click-through detail per §11.4.224); (3) CODE COVERAGE per §11.4.224 (measured %, green $\geq$ target, amber $\geq$ ratchet-floor, red below); (4) SECURITY POSTURE per §11.4.184 SonarQube scanner + supply-chain integrity per §11.4.246 (reproducible + hermetic builds + SLSA Build Level 2 as fleet-wide minimum); (5) DOCUMENTATION completeness per §11.4.257 (fraction of shipped surfaces with the required manual + guide + FAQ trio); (6) DIAGRAM completeness per §11.4.258 (fraction of shipped surfaces with the required diagram set, non-stale); (7) LIVE HEALTH where applicable (uptime / SLO / recent-crash-rate); (8) OPEN DEFECTS per §11.4.15 (Queued + In-progress + Reopened counts, most-reopened highlighted per §11.4.189); (9) SUPPLY-CHAIN integrity level per §11.4.246 (the SLSA Build Level tracked in `docs/security/SLSA_LEVEL.md`, fleet-wide minimum L2); (10) ZERO-SHORTCOMINGS status per §11.4.261 (audit-sweep verdict); (11) MACHINE-EVIDENCE COVERAGE per §11.4.262 (fraction of PASS claims with captured evidence paths); (12) PRODUCTION READINESS gauge (see clause (D)).

(C) MACHINE-DERIVED SOURCE DATA. Every badge's color + value is COMPUTED from a live source of truth, NEVER hand-typed. Every badge PROVENANCE — the exact source path/query/build id — is embedded in the badge tooltip OR a companion `docs/BADGES.md` provenance table. A badge whose provenance is `hand-typed` fails §11.4.259.

(D) PRODUCTION-READINESS GAUGE. A composite badge — GREEN when every clause-(B) badge is GREEN AND no §11.4.185 manual-QA blocker is open AND no §11.4.260 production-readiness criterion is unmet; AMBER when one or more badges are AMBER but no RED and no blocker; RED when any badge is RED or any blocker is open. The gauge is the operator's one-glance answer to "can

we deploy this to production?" and its color is BINDING on the release decision — a RED gauge is a release-blocker (§11.4.108 layer 4 + §11.4.129 huge-blocker family + §11.4.185 manual-QA + §11.4.236 QA-deploy-readiness).

(E) ALWAYS IN SYNC. Badges regenerate on every state change (build completion, test run, defect open/close, doc change, release cut) via §11.4.106 Docs-Chain at the same write-seams that regenerate the rest of the doc surface; a stale badge (color no longer reflects source data) is a §11.4.229 in-sync violation at the reader-facing layer.

(F) BADGE SELF-VALIDATED. The badge-computer is a guard (§11.4.201): golden-GREEN fixture MUST render GREEN, golden-RED fixture MUST render RED, golden-AMBER MUST render AMBER; a badge-computer that PASSES its golden-RED fixture is the §11.4 bluff itself and a release-blocker. Paired §1.1 mutation flips the golden-RED source data and confirms the badge flips.

Honest boundary (§11.4.6). §11.4.259 mandates every applicable badge is PRESENT + accurately derived; it does NOT claim green badges prove the project is bug-free (badges are necessary-not-sufficient, mirroring the §11.4.224 metric-validation family); a project whose badge row is HONESTLY red is compliant with §11.4.259 (the reader gets the truth) even while it fails other release gates. A badge class that GENUINELY does not apply is GRAY-labeled with a reason, never omitted (silence-as-badge is a §11.4.201(6) FALSE-NULL — the reader assumes green because they see nothing red).

Classification: universal (§11.4.17) — badge PALETTE + AT-TOP-OF-README placement + closed color vocabulary + machine-derived + production-readiness gauge are universal; the SPECIFIC badge palette per project is consumer-owned §11.4.35 DATA. Composes §11.4.6 / §11.4.15 / §11.4.27 / §11.4.44 / §11.4.57 / §11.4.86 / §11.4.93 / §11.4.106 / §11.4.108 / §11.4.129 / §11.4.135 / §11.4.169 / §11.4.184 / §11.4.185 / §11.4.189 / §11.4.201 / §11.4.212 / §11.4.223 / §11.4.224 / §11.4.229 / §11.4.236 / §11.4.238 / §11.4.257 / §11.4.258 / §11.4.260 / §11.4.261 / §11.4.262. Propagation gate CM-COVENANT-114-259-PROPAGATION (literal 11.4.259) + recommended gates CM-README-BADGE-ROW-AT-TOP + CM-BADGE-CLOSED-COLOR-VOCABULARY + CM-BADGE-MACHINE-DERIVED-SOURCE + CM-PRODUCTION-READINESS-GAUGE + CM-BADGE-SELF-VALIDATED + paired §1.1 mutations.

Canonical authority: constitution submodule `Constitution.md` §11.4.259. Non-compliance is a release blocker. No escape hatch — no `--badges-optional`, `--hand-type-color`, `--skip-production-gauge`, `--badge-not-derived-from-source`^{2§0}, `--silent-omit-red-badge`, `--badge-in-footer-OK` flag.

§ . . — Cutting-edge enterprise quality + production-deployment readiness: every work product is built for production deployment, maximal stability, zero nasty surprises, and enterprise-grade robustness — never for "good enough" or "will do for now" (BACKGROUND :: REMINDER :: IMPORTANT operator mandate, — —)

Verbatim operator mandate (2026-08-15, Point 6): *"All work we do MUST BE done for the cutting edge enterprise quality! We are doing everything for the deployment to production and maximal stability with no nasty surprises or issues of any kind!"*

Compact summary: every work product a consuming project ships MUST be built for CUTTING-EDGE ENTERPRISE QUALITY targeted at PRODUCTION DEPLOYMENT + MAXIMAL STABILITY + ZERO NASTY SURPRISES — the standing default posture, engaged automatically for every batch, subagent, feature, fix, and change from the first prompt of the session onward (§11.4.126 default-autonomous-loop composed with §11.4.198 always-on defaults), NEVER a per-request opt-in and NEVER downgraded to "good-enough" or "will-do-for-now" or "MVP-first"; every change MUST satisfy the closed PRODUCTION-READINESS invariant set (clause B) before it is considered done, and a project maintains a `docs/PRODUCTION_READINESS.md` document (§11.4.153 sibling) that tracks readiness per component × invariant × evidence-source. §11.4.260 STRENGTHENS §11.4.190 (bleeding-edge enterprise visual quality — WEB surface only) into a UNIVERSAL production-readiness discipline across every surface (web, service, CLI, container, library, docs, orchestration, tests, gates) and BINDS §11.4.5/§11.4.69/§11.4.107 (captured evidence at every claim) + §11.4.108 (four-layer verification) + §11.4.169 (comprehensive test-type coverage) +

§11.4.185 (manual-QA sufficiency) + §11.4.226 (evidence-class-at-closure) + §11.4.235 (build-when-source-proven-correct) + §11.4.236 (QA-deploy-readiness) into one always-on production-readiness posture.

(A) CUTTING-EDGE ≠ NOVELTY-FOR-ITS-OWN-SAKE. "Cutting-edge" means the current best-known engineering approach for the problem class, sourced via §11.4.99 latest-source verification (never memory) + §11.4.150 deep multi-angle research + §11.4.8 external precedent, applied fully — NEVER a shortcut, a stub, a placeholder-with-a-TODO, a `for now` (§11.4.124), a duct-tape workaround, or a "will fix later" (§11.4.197 loss-of-requirements); if the current best-known approach is genuinely unavailable in the project's constraints, that gap is HONESTLY tracked (§11.4.197 item + `[OPEN]` marker per §11.4.223), NEVER silently absorbed as "we did our best."

(B) PRODUCTION-DEPLOYMENT-READINESS INVARIANT SET. Every shipped work product satisfies ALL of: (1) FUNCTIONAL COMPLETENESS — every stated capability wired + verified on a clean target per §11.4.108, no unwired feature, no dead-code path (§11.4.124); (2) OBSERVABILITY — health endpoint, structured logs, per-operation metrics, distributed-trace-ready surface, evidence-capture at every gate (§11.4.85 stress+chaos telemetry, §11.4.128 always-on recording); (3) RESILIENCE — real failure-injection tested per §11.4.85 (chaos: process death, network fault, resource exhaustion, state corruption); recovery is verified never assumed; (4) SECURITY — SonarQube CLI installed + scanned (§11.4.184), no leaked credentials (§11.4.10 pre-store audit), supply-chain integrity at the fleet minimum level (SLSA Build Level 2 per §11.4.246); (5) PERFORMANCE — measured under sustained load per §11.4.85, latency + throughput + resource footprint within declared SLOs; (6) MAINTAINABILITY — code review passed at Fable-xhigh (§11.4.209) with zero findings + zero warnings iterated to GO per §11.4.134; (7) DEPLOYABILITY — rootless container per §11.4.161, deterministic build + reproducible artifact per §11.4.108 layer 2, deployment verified via §11.4.200 target-isolation + §11.4.108 layer 3 runtime-signature; (8) DOCUMENTABILITY — §11.4.257 (manuals/guides/FAQ) + §11.4.258 (diagrams) + §11.4.259 (badges) all GREEN; (9) TESTABILITY — §11.4.27/§11.4.169 seven-type breadth + §11.4.224 coverage floor, every gate paired §1.1 mutation, every regression permanently guarded per §11.4.135; (10) OPERATIONAL RUNBOOK — for every non-trivial failure mode a documented remediation path (§11.4.257 guide class).

(C) ZERO NASTY SURPRISES. No known-broken behavior ships (§11.4.238 QA discovery channel closed), no failing gate is silenced (§11.4.120 seam-placement + §11.4.201 real-condition), no test disabled without a tracked follow-up (§11.4.197), no closed-with-workaround defect masquerading as fixed (§11.4.34 reason-source), no "will fail sometimes but rarely" tolerance (§11.4.50 deterministic consistency + §11.4.85 stress/chaos), no try/catch/swallow silent-failure paths (a caught error either has a captured recovery path or re-raises to a bounded operator decision per §11.4.101/§11.4.66).

(D) STANDING DEFAULT. Composes §11.4.126 (default autonomous-loop mode) + §11.4.198 (always-on default mechanisms) + §11.4.183 (maximum-useful multi-agent + full-constitution-application per track) + §11.4.231 (nano-precision model + effort tier selection); the production-readiness discipline is engaged from the FIRST prompt of the session, never negotiated, never opt-in.

(E) docs/PRODUCTION_READINESS.md — the tracker. Per-component × invariant × evidence-source table (§11.4.153 sibling); revision header per §11.4.44; auto-synced via §11.4.106; reachable from README per §11.4.212; the source of truth for the §11.4.259 production-readiness gauge.

Honest boundary (§11.4.6). §11.4.260 mandates the POSTURE + the INVARIANT SET + the TRACKER; it does NOT claim any project is provably-bug-free (§11.4.118 discovery-pressure boundary remains, §11.4.238 escape triggers a coverage audit not a proof); it does NOT license shipping with unmet invariants — an unmet invariant is HONESTLY red-badged (§11.4.259) and blocks release. "Cutting-edge" is bounded to the project's problem class + published best-known approaches — a genuine research frontier where no best-known exists is HONESTLY marked as such and the project's own approach is documented as an original work contribution (§11.4.8 "NO external solution found — original work" citation).

Classification: universal (§11.4.17). Composes §11.4.5 / §11.4.6 / §11.4.8 / §11.4.10 / §11.4.27 / §11.4.34 / §11.4.44 / §11.4.50 / §11.4.66 / §11.4.69 / §11.4.85 / §11.4.99 / §11.4.101 / §11.4.106 / §11.4.107 / §11.4.108 / §11.4.118 / §11.4.120 / §11.4.124 / §11.4.126 / §11.4.128 / §11.4.134 / §11.4.135 / §11.4.150 / §11.4.153 / §11.4.161 / §11.4.169 / §11.4.183 / §11.4.184 / §11.4.185 / §11.4.190 / §11.4.197 / §11.4.198 / §11.4.200 / §11.4.201 / §11.4.209 / §11.4.212 / §11.4.223 / §11.4.224 / §11.4.226 / §11.4.231 / §11.4.235 / §11.4.236 / §11.4.238 / §11.4.257 / §11.4.258 / §11.4.259 / §11.4.261 / §11.4.262. Propagation gate CM-COVENANT-114-260-

PROPAGATION (literal 11.4.260) + recommended gates CM-PRODUCTION-READINESS-TRACKER-PRESENT + CM-CUTTING-EDGE-POSTURE-ALWAYS-ON + CM-ZERO-NASTY-SURPRISES-AUDIT + paired §1.1 mutations.

Canonical authority: constitution submodule Constitution.md §11.4.260. Non-compliance is a release blocker. No escape hatch — no --mvp-quality-OK , --good-enough , --will-fix-later , --skip-production-readiness-tracker , --production-readiness-optional , --acceptable-flakiness , --for-now-implementation flag.

§ . . – Zero-shortcomings / zero-gaps / zero-weak-spots / zero-danger-zones invariant with a mechanical audit ratchet: at all times the project holds exactly zero of these; every finding is closed OR tracked-with-honest-mitigation, and the audit-run count of open findings is MONOTONE-DECREASING (BACKGROUND :: REMINDER :: IMPORTANT operator mandate, – –)

Verbatim operator mandate (2026-08-15, Point 7): *"There MUST BE ALWAYS exactly zero shortcomings, gaps, weak spots or danger zones in the project in any of the code we have, services or components!"*

Compact summary: every consuming project MUST hold, at ALL TIMES, exactly ZERO shortcomings, gaps, weak spots, or danger zones across every code path, service, component, gate, doc, and configuration — a standing invariant, not a per-release checklist; enforced by a MECHANICAL AUDIT SCRIPT that inventories every known finding class from a closed vocabulary (clause B), runs at every pre-build + release + constitution-pull seam, produces a machine-readable ledger, and is bound to a MONOTONE-DECREASING ratchet (§11.4.135 pattern) so the count of open findings only ever drops or the release refuses; every finding is either FULLY CLOSED with captured evidence (§11.4.5/§11.4.69/§11.4.107 + §11.4.226 evidence-class-at-closure) or explicitly TRACKED as a §11.4.197 item with an honest §11.4.223 provenance marker + a §11.4.101-reversible-safe mitigation. "Zero" is the standing invariant; the ratchet is the mechanism that makes it reachable from a brownfield state without silently absorbing debt. §11.4.261

GENERALISES §11.4.124 (dead-code investigate-before-remove) + §11.4.118 (discovery pressure — enumerated but unratched) + §11.4.238 (QA escape triggers audit — one channel) into a WHOLE-PROJECT zero-finding invariant across the closed set.

(A) CLOSED FINDING VOCABULARY — an audit-actionable enumeration, extensible via named consumer-owned additions per §11.4.35 but never subtractable: (1) SHORTCOMINGS — implementation gaps behind a working contract; (2) GAPS — missing capabilities the docs / schema / API promise; (3) WEAK SPOTS — code paths without adequate guards / error handling / input validation; (4) DANGER ZONES — code / config / operational patterns known to cause serious harm under conditions; (5) TODOs / FIXMEs / "for now" / "temporary" / "will do later" / "placeholder" / `NotImplementedError` / stub-returns / dead code (§11.4.124); (6) SKIPPED tests / disabled gates / commented-out assertions without a tracked `# SKIP-OK: <ticket>` per §11.4.3; (7) BLUFFS — any §11.4/§11.4.1 covenant violation surface (PASS-bluff, FAIL-bluff, echo-observable, source-green-closes-runtime); (8) UNRESOLVED §11.4.197 items past their stated completion window; (9) DIVERGENT / STALE / ORPHAN artifacts per §11.4.233 anti-mess control plane; (10) UN-CATALOGUED anti-patterns identified in code review that lack a §1.1 mutation.

(B) MECHANICAL AUDIT SWEEP — a single executable script (Lava binding: `scripts/audit/zero_findings_sweep.sh` ; consumers supply their path as §11.4.35 DATA) that iterates the closed vocabulary and emits a machine-readable finding ledger `{finding_id, class, file, line, description, status ∈ {open|tracked|closed}, tracker_ref, mitigation_evidence}` — persistent, append-only + snapshot per §11.4.116, git-tracked (§11.4.95) so the ledger IS the SSoT; runs at pre-build + release-tag + constitution-pull; each finding class ships golden-good/golden-bad fixtures per §11.4.107(10) so the audit itself cannot bluff.

(C) MONOTONE-DECREASING RATCHET. The ratchet snapshot records the accepted open-finding count at brownfield adoption (per the §11.4.135 pattern, per-class + total); every subsequent audit run's count MUST be $\leq$ the ratchet snapshot for each class; a run that increases any class's count REFUSES the seam (pre-build gate FAILs, release refuses, commit refuses); the ratchet may DECREASE (a manual snapshot-lower after a finding is closed) and MUST NEVER INCREASE —

the release-tag path never silently absorbs new debt. Brownfield adoption per operator §11.4.66 decision (per §11.4.224(E) adoption-fence pattern), never an invented ratchet.

(D) EVERY FINDING → EITHER CLOSED OR TRACKED (never silently absorbed). CLOSED requires captured evidence at the defect layer per §11.4.226 (runtime > artifact > source), the code / config genuinely repaired, the finding disappears from the audit sweep; TRACKED requires (i) a §11.4.197 workable item with §11.4.148 comprehensive structured description, (ii) an honest §11.4.223 provenance marker ([OPEN: ATM-NNN]) at the finding site, (iii) a reversible-safe mitigation per §11.4.101, (iv) the tracked item's own completion is on the standing autonomous loop's queue (§11.4.87/§11.4.94/§11.4.97/§11.4.103/§11.4.126); (v) a §11.4.135 permanent regression guard where the tracked item's exposure could cause harm before closure. Silent absorption ("we know this is broken but shipping anyway") is a §11.4.261 violation of §11.4 PASS-bluff severity — the ledger IS the honesty seam.

(E) HONEST "ZERO" — never the fake kind. "Exactly zero" means the CLOSED-VOCABULARY-audit-run finds zero unclosed-untracked entries; it does NOT claim omniscience (§11.4.118 discovery-pressure boundary — un-catalogued finding classes stay possible, and the vocabulary is EXTENSIBLE for exactly that reason); a §11.4.238 escape triggers the class-extension of the audit + a ratchet-snapshot bump-DOWN (never up) once the escaped class is added and cleared. Honesty at the "zero" claim requires the vocabulary is genuinely covered — a project narrowing the vocabulary to declare "zero" is a §11.4.6 fabrication and a release-blocker.

Honest boundary (§11.4.6). §11.4.261 mandates the STANDING INVARIANT + the AUDIT + the RATCHET + the TRACKED-vs-CLOSED discipline; it does NOT prove the project is provably-flawless, and does NOT license a "zero" claim over a narrowed vocabulary. Composes with, never substitutes for, §11.4.118 (§11.4.261 catches the CATALOGUED — discovery still needed for the un-catalogued), §11.4.238 (an escape triggers vocabulary extension), §11.4.226 (evidence class at closure decides whether a finding is genuinely closed), §11.4.135 (regression guards for closed items keep them closed).

Classification: universal (§11.4.17). Composes §11.4.3 / §11.4.6 / §11.4.15 / §11.4.16 / §11.4.34 / §11.4.44 / §11.4.66 / §11.4.87 / §11.4.94 / §11.4.95 / §11.4.97 / §11.4.101 / §11.4.103 / §11.4.107(10) / §11.4.113 / §11.4.115 / §11.4.116 /

§11.4.118 / §11.4.124 / §11.4.126 / §11.4.135 / §11.4.148 / §11.4.161 / §11.4.197 / §11.4.201 / §11.4.223 / §11.4.224(E) / §11.4.226 / §11.4.233 / §11.4.238.

Propagation gate CM-COVENANT-114-261-PROPAGATION (literal 11.4.261) + recommended gates CM-ZERO-FINDINGS-AUDIT-SWEEP + CM-ZERO-FINDINGS-MONOTONE-RATCHET + CM-EVERY-FINDING-CLOSED-OR-TRACKED + paired §1.1 mutations.

Canonical authority: constitution submodule Constitution.md §11.4.261. Non-compliance is a release blocker. No escape hatch — no --skip-findings-audit , --allow-todo-without-tracker , --ratchet-may-increase , --silently-absorb-finding , --narrow-vocabulary-to-declare-zero , --for-now-in-production flag.

§ . . — Machine-created evidence at every gate: every claim of "works" / "passes" / "verified" cites a^{0.2.6} captured, machine-derived, machine-verifiable evidence artifact produced by the gate — no operator eyeballing, no narrative-only PASS (BACKGROUND :: REMINDER :: IMPORTANT operator mandate, — —)

Verbatim operator mandate (2026-08-15, Point 8): *"There MUST BE no bluff of any kind and in any form anywhere! Everything created MUST BE fully covered with confirmation, validation and verification gates which will assert machine created evidence so there is no prediction or guessing about the success and production ready state of the System!"*

Compact summary: every claim by the System that something WORKS / PASSES / IS VERIFIED / IS PRODUCTION-READY MUST be backed by MACHINE-CREATED, MACHINE-VERIFIABLE evidence — a file / verdict / measurement produced by the gate itself, cited by path + content-hash + timestamp, with a schema an independent verifier can parse (§11.4.226 evidence-class-at-closure applied to EVERY gate, not only closure seams); operator eyeballing, human narrative, "I checked it works", "looks fine", "no error was reported", and every other prediction / guess / assumption about success is FORBIDDEN at the evidence layer (§11.4/§11.4.1/§11.4.6). §11.4.262 STRENGTHENS the §11.4 anti-bluff covenant preamble from "captured evidence"

to MACHINE-CREATED evidence (a captured screenshot an operator interpreted is captured but not machine-verified — §11.4.107(10) golden-good/golden-bad self-validating analyzer applies UNIVERSALLY, not only to AV playback), and BINDS §11.4.5/§11.4.69/§11.4.107/§11.4.108/§11.4.115(F)/§11.4.201/§11.4.226 into one universal machine-evidence-at-every-gate discipline. Prediction, guessing, and assumption about success are the exact anti-patterns this anchor forbids.

(A) THREE-GATE COVERAGE. Every work product passes THREE gate classes each producing machine-created evidence: (1) CONFIRMATION — the change lands correctly (source review, static analysis, format/lint, dependency clash check per §11.4.110); (2) VALIDATION — the change behaves as designed at each of the §11.4.108 four layers (SOURCE / ARTIFACT / RUNTIME / USER-VISIBLE), each producing its own machine-created verdict + evidence artifact; (3) VERIFICATION — the change satisfies its acceptance contract from the user's / operator's / system's point of view (§11.4.185 manual-QA is the human sufficiency gate; §11.4.236 QA-deploy-readiness is the machine sufficiency gate — each with its own captured evidence).

(B) MACHINE-CREATED — the four properties. Every evidence artifact is (i) PRODUCED BY THE GATE, not authored by the tester (a screenshot the harness took, not one the operator narrated); (ii) MACHINE-VERIFIABLE — a schema-parseable file (JSON verdict, JSONL event stream, structured log, image with an OCR/vision oracle per §11.4.107(11)/§11.4.117, audio with RMS/loudness analysis per §11.4.107(11)) that an INDEPENDENT verifier can re-check without asking the tester; (iii) CONTENT-ADDRESSED — cited by path + sha256 + timestamp (§11.4.207 content-addressed Merkle) so tampering is detectable; (iv) SEMANTIC-CLASS-MATCHED — the evidence class matches the claim layer (§11.4.226 wrong-layer echo forbidden — a source grep transcript can never satisfy a "pixel renders green" claim).

(C) NO NARRATIVE-ONLY PASS. A PASS whose evidence is a paragraph of natural language authored by an operator or a subagent ("I ran the tests and they all passed", "looks good to me", "manually confirmed") is a §11.4.262 violation — the narrative may accompany the machine evidence but MUST NEVER substitute for it. §11.4.185 manual-QA remains a human-sufficiency gate whose OWN evidence is machine-created (a QA session recording per §11.4.128, a QA checklist file with per-item PASS/FAIL + evidence-path per §11.4.185, a §11.4.238 escape-log if the human found something automation missed).

(D) SELF-VALIDATED EVIDENCE ANALYZERS. Every evidence analyzer (the parser that reads the machine-created artifact and emits a verdict) is itself golden-good + golden-bad + negative-control validated per §11.4.107(10); an analyzer that PASSES its golden-bad fixture is the bluff and a release-blocker (the §11.4.201 self-validation discipline applied to evidence tooling). Every claim of "works" cites the analyzer's verdict + the raw evidence, so an independent re-run can reproduce the verdict.

(E) FOUR-LAYER BINDING WITH §11.4.108. At each of the four verification layers (SOURCE / ARTIFACT / RUNTIME-ON-CLEAN-TARGET / USER-VISIBLE), the layer's PASS cites MACHINE-CREATED evidence PRODUCED AT THAT LAYER (source: static-analysis JSON + review-artifact; artifact: build-log + fingerprint-checksum + byte-check of the change's landed bytes; runtime: on-target verdict + captured runtime signature + telemetry; user-visible: OCR/vision oracle output or sink-side probe report per §11.4.13/§11.4.69/§11.4.107). A PASS at a higher layer that cites evidence from a lower layer alone is a §11.4.226 wrong-layer violation — the machine-evidence chain runs THROUGH the layers, never SKIPS.

(F) STANDING DEFAULT (§11.4.126/§11.4.198). Machine-evidence-at-every-gate is the STANDING DEFAULT posture engaged from the first prompt of the session; every subagent dispatch, every review, every test run, every gate emits its evidence artifact into the project's evidence store (§11.4.207 content-addressed) so the full audit trail is machine-searchable at any time by any operator / subagent / reviewer. Prediction, guessing, and assumption about success are forbidden at the reporting layer as well as the gate layer — a status report that predicts "should pass" without citing a run's evidence is a §11.4/§11.4.1 bluff at the reporting layer.

Honest boundary (§11.4.6). §11.4.262 mandates every gate's PASS is BACKED by machine-created evidence + the analyzer is self-validated; it does NOT claim machine evidence proves human satisfaction (§11.4.185 remains, with its own machine-evidence requirement per (C)); it does NOT license shipping when the machine evidence disagrees with reality (a bug the automation missed is a §11.4.238 escape triggering a coverage audit + a vocabulary extension per §11.4.261, not a "the machine said green so we're good" defense). "Machine-created" is bounded to the artifact — the DECISION about whether a gate passed remains policy the humans authored (§11.4.6 no-guessing extends to policy calibration, not to the gate's mechanical execution).

Classification: universal (§11.4.17). Composes §11.4/§11.4.1 (anti-bluff covenant) / §11.4.5 / §11.4.6 / §11.4.13 / §11.4.15 / §11.4.20 / §11.4.27 / §11.4.50 / §11.4.66 / §11.4.69 / §11.4.85 / §11.4.98 / §11.4.107(10)(11)(12) / §11.4.108 / §11.4.110 / §11.4.115(F) / §11.4.116 / §11.4.117 / §11.4.126 / §11.4.128 / §11.4.135 / §11.4.169 / §11.4.185 / §11.4.198 / §11.4.201 / §11.4.207 / §11.4.226 / §11.4.236 / §11.4.238 / §11.4.259 / §11.4.260 / §11.4.261. Propagation gate CM-COVENANT-114-262-PROPAGATION (literal 11.4.262) + recommended gates CM-MACHINE-EVIDENCE-AT-EVERY-GATE + CM-EVIDENCE-ANALYZER-SELF-VALIDATED + CM-NO-NARRATIVE-ONLY-PASS + CM-EVIDENCE-LAYER-MATCH-108 + paired §1.1 mutations.

Canonical authority: constitution submodule Constitution.md §11.4.262. Non-compliance is a release blocker. No escape hatch — no --narrative-pass-OK , --operator-eyeball , --should-pass-suffices , --skip-evidence-analyzer-validation , --evidence-optional , --pass-without-machine-artifact , --predict-success flag.

§ . . — Process-group signal-safety mandate: NEVER signal pgid ≤ , NEVER trust a mock-derived pid/pgid — validate as int > before every killpg / kill(-pid, sig) / pkill / killall call (BOB- forensic anchor, - -)

Forensic anchor — verbatim operator mandate (2026-08-19):

"CRITICAL: Fix MUST BE systematic so any other projects do not have same issue!"

Forensic case study (BOB-126, boba project, 2026-08-19). SEVEN forced-logout incidents on the operator's workstation over ~48 hours (BOB-116/120/123/124/125/126) — the operator's entire graphical session, tmux, SSH, browsers, and Claude Code process were repeatedly SIGKILLed. Kernel audit rules installed 2026-08-19 15:56 finally captured the initiator on incident #7:

```
audit[399861]: SYSCALL syscall=62 a0=ffffffff a1=9
pid=399861 auid=1000 comm="pytest" exe="/usr/bin/python3.14"
key="sigkill_investigation"
```

A `pytest` process called `kill(-1, SIGKILL)`. Root cause: a Python test created an `AsyncMock()` subprocess object without explicitly setting `mock.pid` as an int. The production code path called `os.killpg(os.getpgid(proc.pid), signal.SIGKILL)`. Python's `MagicMock.__int__` defaults to `1`, so:

1. `proc.pid` → `MagicMock`
2. Python coerces to `int` for `os.getpgid` → `int(MagicMock()) == 1` (documented default)
3. `os.getpgid(1)` returns `1` (init's process group)
4. `os.killpg(1, SIGKILL)` → glibc → `kill(-1, SIGKILL)`
5. `kill(-1, sig)` semantics = "signal every process the caller may signal, except pid 1 and self" → under UID 1000 that means the entire user session (systemd `user@1000` manager, gnome-shell, tmux, ssh, browsers, Claude Code)

The `contextlib.suppress(Exception)` around the call swallowed nothing because the syscall SUCCEEDED before any exception path. The defect existed for ~4 months (since 2026-04-24) before audit rules exposed it.

The mandate (ALL hold):

(A) NEVER signal $\text{pgid} \leq 1$. No project, no language binding, no code path may call `os.killpg(pgid, sig)` / `kill(-pid, sig)` / `killpg(2)` (POSIX) / `pkill -<pgid>` / `subprocess.killpg` / `Process.killpg` where `pgid` (or `pid` in the negative-pid form of `kill(2)`) is ≤ 1 . On Linux, `killpg(1, sig)` and `kill(-1, sig)` are equivalent and BOTH signal every process the caller has permission to signal — under a non-root UID, that includes the user's session manager, GUI, shell, and every daemon in the same UID scope. This is the disaster syscall. Guarding with a positive integer test is trivial; failing to guard is catastrophic. On other Unixes (BSD family, macOS) the semantics of `kill(-1, sig)` is broadly the same "signal all processes the caller may signal"; the mandate is portable.

(B) VALIDATE `pid + pgid` as `int > 1` BEFORE every process-group signal call. The universal invariant:

```
IF calling killpg(pgid, sig) OR kill(-pid, sig) OR any process-group
signal:
    ASSERT isinstance(pgid, int)           # never a mock, never a string,
never None
    ASSERT pgid > 1                       # never init's pgid, never 0 (self-
group), never negative
    (same for pid when the negative-pid kill(2) form is used)
```

Language-specific enforcement:

- **Python:** guard with `if isinstance(pid, int) and pid > 1: pgid = os.getpgid(pid); if isinstance(pgid, int) and pgid > 1: os.killpg(pgid, sig) .`
- **Go:** `syscall.Kill(-pgid, sig)` — guard `pgid > 1` before the call; test doubles must return `int(pgid) > 1` .
- **Rust:** `nix::sys::signal::killpg` — accept only `Pid::from_raw(n)` with `n > 1` , and reject the sentinel meaning "current process group".
- **Bash / shell:** `kill -<pgid> <sig> / pkill -g <pgid>` — validate `[[ "$pgid" -gt 1 ]]` before the call.
- **C:** `killpg(pgid, sig)` and `kill(-pid, sig)` — same integer guard.

(C) TESTS mocking subprocess/proc objects MUST explicitly set `mock.pid` (or equivalent) as `int`. A subprocess mock with an unset `pid` field is BOB-126-shaped by construction: the language's mock library (Python's `MagicMock` , Go testify's `Mock` , JS Jest's `jest.fn()` , etc.) provides a default coercion that can silently become a low integer. Every test that creates a fake subprocess object MUST:

1. Set `mock.pid = <large int, e.g. 12345>` explicitly — never rely on the mock library's default.
2. Patch the process-group signaling call itself (e.g. `patch.object(module.os, 'killpg')`) as belt-and-suspenders so a future regression of the production code's guard cannot re-open the disaster.
3. Assert the mocked `killpg` was called with `pgid > 1` (paired §1.1 mutation catches a regression that lets `pgid=1` through).

(D) DEFENSE-IN-DEPTH (§11.4.107(10) analyzer discipline applied). The guard MUST be present in BOTH the production code AND the test. Removing either layer alone must break at least one test. This is the four-layer coverage per §11.4.4(b) applied to the signal-safety class:

- Layer 1 (pre-build gate): grep pattern for `killpg\(.*\)` / `kill\(-.*\)` calls that are NOT preceded by the guard — surface as review findings.
- Layer 2 (post-build gate): CodeGraph query for callers of `os.killpg` / `syscall.Kill(-*)` / equivalent, cross-referenced against guard presence.
- Layer 3 (runtime): kernel audit rules on SIGKILL syscalls with `a0=0xffffffff` (`kill(-1)`) or `a0=0x00000001` in the `pidfd_send_signal` case — a live host-safety canary.
- Layer 4 (paired §1.1 mutation): a regression test that DELIBERATELY invokes the `killpg` code path with a mock pid AND asserts `killpg` was not called with `pgid ≤ 1`. Reverts of the pid guard MUST make this test fail.

(E) HONEST BOUNDARY (§11.4.6). §11.4.263 hardens the process-group signal APIs specifically. It does NOT prohibit legitimate use of `killpg` for well-scoped process groups (`start_new_session=True` subprocesses where the child owns its group). It does NOT extend to signals other than SIGKILL where the blast radius is small (SIGCHLD, SIGUSR1, etc. — though the same guard is cheap and recommended). It does NOT prevent kernel OOM-killer or systemd-oomd from killing the user (those are separate host-safety concerns — §12.6 memory ceiling + §11.4.225 quota telemetry apply).

(F) STANDING DEFAULT. Every project inheriting this constitution treats §11.4.263 as always-on from the first prompt. Every code review (§11.4.125/§11.4.142/ §11.4.194) MUST scan for process-group signal calls and verify the guard. Every subagent dispatched to write subprocess-cleanup code MUST include the guard in its output. Every test-writing subagent MUST include the `mock.pid = <int>` discipline.

Classification: universal (§11.4.17). Composes §11.4.1 (FAIL-bluff — swallowed exception around a successful catastrophic syscall is a bluff at the error-handling layer) / §11.4.4(b) (four-layer coverage) / §11.4.5 / §11.4.6 / §11.4.85 (chaos-test host safety) / §11.4.107(10) (self-validated analyzer) / §11.4.115 (RED-first with the actual defect condition) / §11.4.126 (standing default) / §11.4.201 (guard asserts real condition — `pgid > 1`) / §11.4.225 (host-safety orthogonal axis) / §11.4.226 (real

captured evidence). Propagation gate `CM-COVENANT-114-263-PROPAGATION` (literal `11.4.263`) + recommended gates `CM-KILLPG-PGID-GUARD` (production code — every `killpg/kill(-pid)` call site has the `pid>1 && pgid>1` guard) + `CM-TEST-MOCK-PID-EXPLICIT-INT` (test code — every subprocess mock has `mock.pid` explicitly set as int + `killpg` patched) + paired §1.1 mutations (remove the guard from production → the regression test fails; unset `mock.pid` in the test → assertion catches `pgid<=1` attempt).

Canonical authority: constitution submodule `Constitution.md` §11.4.263. Non-compliance is a release blocker regardless of context. No escape hatch — no `--skip-killpg-guard`, `--mock-pid-optional`, `--killpg-1-is-ok`, `--kill-minus-1-permitted`, `--suppress-catches-syscall`, `--test-may-not-explicit-pid` flag exists.