

Helix Universal Constitution

Field	Value
Revision	27
Created	2026-05-14
Last modified	2026-06-13T00:00:00Z
Status	active
Status summary	Added §11.4.152 — Crashlytics-recorded-data continuous monitoring + systematic-debug + regression-test-coverage mandate (User mandate 2026-06-13): every project with Firebase Crashlytics enabled/wired MUST continuously monitor ALL four recorded surfaces — fatal crashes, ANRs, performance traces, AND non-fatals — systematic-debug each (reproduce-before-fix §11.4.102/.115), fix/improve, and cover every closure with a permanent §11.4.135 regression guard (RED-on-broken→GREEN-on-fixed) + a closure log citing the console issue id/URL + the validation/verification test paths; on a regular cadence (the §11.4.47 five-trigger set), no false results, no bluff — a console-resolved Crashlytics issue with no falsifiable regression test is FORBIDDEN (the silent recurrence vector). STRENGTHENS+COMPLETES §11.4.47 (review/dedup finds it; §11.4.152 fixes it + proves it stays fixed). Project-level reference instantiation = the Lava client’s §6.O (Crashlytics-Resolved Issue Coverage) + §6.AC (Comprehensive Non-Fatality Telemetry). Composes §11.4/.1/.5/.6/.9/.34/.40/.43/.47/.69/.90/.102/.107/.108/.115/.118/.123/.130/.135/.138/.152/§1.1. Propagation gate CM-COVENANT-114-152-PROPAGATION + recommended gate CM-CRASHLYTICS-ISSUE-FULLY-COVERED (gate-code = separate work item). Classification: universal (§11.4.17). CLAUDE.md + AGENTS.md + QWEN.md mirrors updated in lockstep. Prior round: Added §11.4.151 — Project-prefixed release-tag/version-naming mandate (User mandate 2026-06-12): every release tag AND version name on the main repo AND every owned submodule MUST be prefixed <PREFIX>-<version> (e.g. helix_translate-1.0.0-dev-0.0.1); prefix resolution order is (1) HELIX_RELEASE_PREFIX from .env (authoritative, git-ignored §11.4.30, documented in tracked .env.example §11.4.77) else (2) the lowercased snake_case project-root dir name §11.4.29; the SAME prefix across main + all owned submodules in one release makes it greppable across every repo; version codes increment monotonically. Compose §2.1/§11.4.28/.29/.30/.40/.77/.113/.126. Propagation gate CM-COVENANT-114-151-PROPAGATION + recommended gate CM-RELEASE-PREFIX-NAMING (gate-code = separate work item). Classification: universal (§11.4.17). CLAUDE.md + AGENTS.md + QWEN.md mirrors updated in lockstep. Prior round: Added §11.4.148 + §11.4.149 — workable-item integrity + testing-diary family (User mandate 2026-06-10). §11.4.148 Workable-item integrity (status+type+id) + comprehensive structured description + bidirectional external-tracker sync + BLOCKED unblock-choices: BINDS + STRENGTHENS §11.4.15/.16/.21/.54/.91/.93/.95/.106 into ONE integrity contract spanning DB [?] docs [?] external tracker — (D1) no item without a valid status + valid type + stable id on ALL three surfaces; (D2) comprehensive structured description (what manifest/repro/acceptance) in DB + docs + tracker, never a stub; (D3) BLOCKED items carry WHY + enumerated unblock CHOICES (tightens §11.4.21, mirrors §11.4.66 2-4-option shape; Blocked/BLOCKED = documented alias of canonical Operator-blocked)

Field	Value
	(D4) regular never-missed bidirectional DB[?]docs[?]tracker sync, §11.4.86 fingerprint-gated + §11.4.106 docs_chain-bound; (D5) generic external-tracker push carries statuses (collapsed per §11.4.33/.112)/types/assignee (§11.4.104, default-from-env-var, never logged §11.4.10)/sub-tasks, idempotent dry-run-then-real match-by-stable-key, sink-side created/updated/failed proof §11.4.69, machinery project-agnostic §11.4.28. Composes §11.4.15/.16/.21/.33/.34/.54/.66/.86/.91/.93/.95/.104/.106/.112/.123/.10/.28/.69/.6/§1.1. Propagation gate CM-COVENANT-114-148-PROPAGATION + recommended gates CM-ITEM-INTEGRITY-STATUS-TYPE-ID/CM-ITEM-COMPREHENSIVE-DESCRIPTION/CM-BLOCKED-UNBLOCK-CHOICES/CM-TRACKER-SYNC-IDEMPOTENT (gate-code = separate work item). §11.4.149 Per-workable-item testing-diary: every item carries an append-only TESTING DIARY (date/time + tested-by + result + in-depth observations + action taken+why), distinct from §11.4.93 item_history (state) + §11.4.55 Reopens (cycles) — records TEST EXECUTIONS — (a) test_diary table in the §11.4.93/.95 SQLite SSOf with a schema constraint making a PASS-without-evidence-path impossible §11.4.69; (b) in-depth diary + derived at-a-glance summary VIEW feeding §11.4.132 risk-ordering; (c) four-format per-item Diary/Diary_Summary exports §11.4.65 + §11.4.86-fingerprinted §11.4.106 docs_chain-bound (never missed); (d) external-tracker SUB-TASK model with {TODO In-progress Completed} lifecycle (item description stays the item per §11.4.148 D2; diary = children), §11.4.148 D5 idempotent push; (e) minimal-LLM deterministic tooling (prose authored by the tester), 100% test-covered §11.4.27 (unit + integration-against-real-tracker honest §11.4.3 SKIP-when-token-absent + export + HelixQA Challenge + paired §1.1). Honest boundary §11.4.6 — guarantees a complete auditable test-history, NOT that any single verdict is correct. Composes §11.4.6/.27/.45/.50/.55/.65/.69/.86/.93/.95/.106/.107/.123/.132/.148/.10/.13/.28/.33/.112/§1.1. Propagation gate CM-COVENANT-114-149-PROPAGATION + recommended gates CM-TEST-DIARY-SYNC/CM-DIARY-PASS-REQUIRES-EVIDENCE (gate-code = separate work item). Both Classification: universal (§11.4.17); project-specific dual per-display brightness/auto-dim/font-size defaults landed ONLY in the ATMOSphere consumer layer (RK3588/Presenter), NOT here. CLAUDE.md + AGENTS.md + QWEN.md mirrors updated in lockstep. Prior round: Added §11.4.147 — Crashed-agent respawn-until-complete + no-work-loss registry mandate (User mandate 2026-06-10): every dispatched agent/subagent MUST be tracked through its full lifecycle so a crash NEVER loses/forgets/corrupts its work (a crashed agent is NOT a completed agent, §11.4.6 — abnormal termination is positive evidence the unit is OPEN) — (a) durable append-only agent REGISTRY (stable id + §11.4.58 file-scope + §11.4.5/.69 evidence dir + status CLOSED SET {dispatched in-flight crashed respawned complete}) on the §11.4.116 sync substrate, SINGLE SOURCE OF TRUTH for “is this work owed?”; (b) mechanical CRASH-DETECTION (transient rate-limit/socket-close/any non-completion exit) → flip to crashed + keep OPEN → RESPAWN until complete, autonomously per §11.4.101 using ALREADY-DEFINED backoff (never invented §11.4.6), non-transient terminal error investigated per §11.4.102; (c) PARTIAL-STATE preserve → §11.4.84-quiescence-check → resume-idempotently OR clean-restart (§9.2 backup), per-agent git worktree isolation; (d) “crash ≠ done” — the §11.4.87/.94/.97/.126 loop done-condition MUST NOT read satisfied while any entry is non-complete, a complete without landed evidence is a §11.4 PASS-bluff. Forensic anchor (FACT, this session): 5 subagents killed at once by rate-limit, 2 by socket-close, one left PARTIAL dual-sliders edits — re-dispatched by pure conductor vigilance with NO mechanical guarantee; agent side analogue of §11.4.144. Composes §11.4.6/.58/.84/.87/.94/.97/.101/.116/.102/.108/.125/.142/.40/.128/.144/§9.2/§1.1. Propagation gate CM-COVENANT-114-147-PROPAGATION + recommended gate CM-CRASHED-AGENT-RESPAWN-TRACKED (gate-code = separate work item). Classification: universal (§11.4.17). CLAUDE.md + AGENTS.md + QWEN.md mirrors updated in

Field	Value
	lockstep. Prior round: Added §11.4.145 — Independent multi-angle impact-research per change (User mandate 2026-06-10): for EVERY fix/change/new-feature INDEPENDENT impact-research agents (subagent-driven §11.4.70/.20, structurally separate from the author — adversarial “refute-the-change” stance defeating the documented LLM confirmation-bias failure mode) MUST research the change AND its connected/dependent features across a CLOSED SET OF EIGHT ANGLES — (1) correctness/logic, (2) regression (what existing working feature could break — call-graph impact), (3) latent/dangerous-code THE RECENTS CLASS (runtime-only failure, interface/ABI/contract mismatch, concurrency/data-race, resource lifecycle), (4) security, (5) performance, (6) host/data/target-hardware safety §12/§9/.133, (7) cross-feature interaction, (8) business-logic conformance; forensic anchor (FACT) the recents/3-button-nav path shipped a latent AIDL-contract mismatch that PASSED code review yet ANR’d at RUNTIME ONLY (the §11.4.108 SOURCE→ARTIFACT→RUNTIME→USER-VISIBLE gap) → §11.4.145 shifts blast radius discovery LEFT; each angle names the TOOL(S) + obtains the required DEPTH (diff + full unit + declared contract + spec + every connected feature, NEVER diff-lines alone) + emits captured-evidence (§11.4.5/.69, “looks fine” forbidden §11.4.1), N/A angle = NOT_APPLICABLE:<reason> §11.4.6 never silently skipped; output = a per-change impact-research REPORT (one section/angle: tool+evidence+verdict) + a single GO/NO-GO BLOCKING acceptance/commit/build on ANY unmitigated risk, iterating CLEAN GO §11.4.134 before the §11.4.125 gate; honest boundary §11.4.6 — does NOT replace §11.4.108/.40, the FIRST of MULTIPLE STRONG LAYERS. Composes §11.4.92/.125/.108/.110/.134/.78/.79/.107/.85/.133/§12/§9/.99/.6/§1.1. Propagation gate CM-COVENANT-114-145-PROPAGATION + recommended gate CM-IMPACT-RESEARCH-PER-CHANGE (gate-code = separate work item). Classification: universal (§11.4.17). CLAUDE.md + AGENTS.md + QWEN.md mirrors updated in lockstep. Added §11.4.146 — Reproduce-first test + same-test-confirms-fix + mandatory extend-to-all-cases workflow (User mandate 2026-06-10): every reported problem MUST be handled by a NAMED three-step test workflow (§11.4.146 NAMES + BINDS existing disciplines, adds ONLY two emphases) — STEP 1 REPRODUCE-FIRST+INVESTIGATE: author the §11.4.43 RED test as a §11.4.115 RED-baseline-on-the-broken-artifact (capture defect-present physical evidence §11.4.5/.69/.107) + NEW EMPHASIS D1 the same RED test is ALSO a deliberate investigation instrument §11.4.102(A) gathering ADDITIONAL forensic data characterising the defect (triggers/boundaries/scope/adjacent failure modes) feeding BOTH the fix design AND the STEP-extend scope; STEP 2 SAME-TEST-CONFIRMS-FIX: the SAME source (§11.4.115 polarity RED_MODE=1→0) confirms ABSENT on a clean target §11.4.108/.139, validate-fix-first §11.4.130, deterministic §11.4.50, no separate happy-path substitute; STEP 3 EXTEND-TO-ALL-CASES (NEW EMPHASIS D2, mandatory per-fix): FAN OUT across the full case-space of the SAME functionality — all flows + valid/invalid/boundary (§11.4.85) + concurrent (§11.4.85) + chaos (§11.4.85) + topology (§11.4.3) — enumerated per §11.4.118 (never “no other issues found”), each user-visible case physical-evidence §11.4.123 + registered into the §11.4.135 standing suite, a new fan-out defect triggers §11.4.4 + re-enters STEP 1; NOT deferred to a release-cycle sweep, NOT reduced to a single guard. Composes §11.4.43/.115/.102/.130/.108/.139/.50/.85/.118/.135/.3/.123/.5/.69/.107/.138/.4/§107/§1.1. Propagation gate CM-COVENANT-114-146-PROPAGATION + recommended gate CM-REPRODUCE-FIRST-THEN-EXTEND (gate-code = separate work item). Classification: universal (§11.4.17). CLAUDE.md + AGENTS.md + QWEN.md mirrors updated in lockstep. Prior round: Added §11.4.144 — Tracked/recorded-device availability-following mandate (User mandate 2026-06-10): every device the project is tracking / following / recording (test / debug / manual-testing, across every reachable transport — USB / wireless ADB / SSH / serial / network introspection API) MUST be availability-

Field	Value
	FOLLOWED — its connection state continuously monitored, any drop handled, never silently abandoned and never presented as a continuous recording; the §11.4.128 always-on recorder KNOWS a tracked device is absent (its loop guards on the reachable state) but, lacking a following discipline, merely spins idle (no data captured, no offline event no resume, no escalation), so the corpus shows a continuous timeline with a silent hole — a §11.4 PASS-bluff at the recording-integrity layer. On a drop the system MUST automatically (a) DETECT + log an honest offline event (§11.4.6 fabricated-continuity bluff; §11.4.128 always-on gap), (b) WAIT using the project’s ALREADY-DEFINED reconnection timings — never invented numbers (§11.4.6), the SAME grace/reconnect/poll budgets the recovery path already uses, (c) RE-ATTACH + log an honest online/resume event the moment the device returns, (d) ESCALATE to the sanctioned device-recovery path (per §11.4.69 feature class device_recovery) if the device does not return within the defined timeout — through the sanctioned, authorization-gated recovery entry point ONLY, never bypassing its gate, never performing a destructive recovery (power-cycle) autonomously without that authorization (§11.4.21/.101 — high-blast-radius recovery is gated; while blocked keep following + log the blocked-escalation honestly). Composes §11.4.128/.69/.6/.14/.21/.101. Propagation gate CM-COVENANT-114-144-PROPAGATION + recommended gate CM-DEVICE-AVAILABILITY-FOLLOWED (gate-code separate work item). Classification: universal (§11.4.17). CLAUDE.md + AGENTS.md + QWEN.md mirrors updated in lockstep. Prior round: Added §11.4.143 — Real-user-journey mandate for video-streaming-app full-automation tests (User mandate 2026-06-10, verbatim “All video streaming apps ... require to choose some title and to press proper UI button to start or resume playing! Proper UI interaction to play exact show with proper content and subtitles is MANDATORY! Without it we just eventually play on 2nd display sample, and that’s it mostly! THIS MUST BE ADDED as MANDATORY RULE regarding testing of any video streaming app in general with full automation tests! Universal in root constitution + ATMOSphere project extensions.”): any full-automation test asserting a video player / streaming app plays content MUST drive the REAL end-user journey through the app’s OWN UI — launch → BROWSE the actual catalog → choose a SPECIFIC title → press the real Play/Resume button → confirm THAT chosen content plays with its correct subtitles on the intended routing target — NEVER a sample / demo / loop-clip / deep-link / am start -a VIEW / synthetic shortcut (those validate ROUTING only, not real chosen-content playback → §11.4 PASS-bluff at the user-journey layer, the operator’s “we just eventually play on 2nd display sample” gap); non-introspectable streaming UIs use the §11.4.117 CV/OCR pixel oracle; login via the §11.4.10 credential single-source; chosen-content confirmation via §11.4.107 liveness + §11.4.136 real-content + §11.4.137 subtitle correctness; honest operator_attended SKIP-with-reason per §11.4.52/.3 ONLY where autonomous is genuinely infeasible (hard login / geo-block / secure-surface blanking) — NEVER a faked routing-only PASS. Composes §11.4.48/.107/.117/.136/.137/.52/.3/§107/§1.1. Propagation gate CM-COVENANT-114-143-PROPAGATION + recommended gate CM-VIDEO-REAL-JOURNEY-TEST (gate-code = separate work item). Classification: universal (§11.4.17). CLAUDE.md + AGENTS.md + QWEN.md mirrors updated in lockstep. Prior round: Added §11.4.142 — Universal code-review mandate — every change reviewed, always, no exception (User mandate 2026-06-09, verbatim “ALL changes we do MUST pass through the code review step!!! ALWAYS!!!”): EVERY change to ANY governed repository — without exception (source, fixes, tests, gates, meta-test mutation docs, doc-tooling, build/CI scripts, config, governance files, conductor main-stream edit sub-agent output, refactors, single-line edits) — MUST pass an INDEPENDENT code-review step BEFORE acceptance/commit/build; the ABSOLUTE form of §11.4.125 (strips every implicit scoping — no “just a doc edit”, no “just a one-liner”, no “self-review (§11.4.92) suffices”, no “trivial change” carve-out); independence is load-bearing (reviewer structurally separated from the author — dedicated review agent subagent-

Field	Value
	driven per §11.4.70/.20 or distinct human, NEVER author re-reading own work; §11.4.9 self-review PRECEDES never satisfies); the review is itself anti-bluff (§11.4/.1 — a rubber-stamp is not a review; conclusions captured-evidence per §11.4.5/.69) and iterate to a clean GO per §11.4.134 (any finding/warning re-arms; acceptance on ZERO finding + ZERO warnings); honest boundary §11.4.6 — does NOT replace §11.4.108 runtime-signature verification nor §11.4.40 full-suite retest, it is the FIRST of MULTIPLE STRONG LAYERS every change crosses. Composes §11.4.1/.4/.5/.6/.20/.40/.69/.70/.92/.108/.110/.125/.134/§107/§1.1. Propagation gate CM-COVENANT-114-142-PROPAGATION + recommended gate CM-EVERY-CHANGE-REVIEWE (gate-code = separate work item). Classification: universal (§11.4.17). CLAUDE.md + AGENTS.md + QWEN.md mirrors updated in lockstep. Prior round: Added §11.4.125 Code-review-agent gate before pre-build + main build (mandatory multi-layer review) (User mandate 2026-06-04, verbatim “After all fixes/changes/implementations are done BEFORE running pre-build tests and the main build, dispatch code-review agent(s) that analyze all work done + all existing data/facts + the existing codebase + current git history to determine quality, safety, and whether the fixes/changes will REALLY work; they MUST validate and verify that every test covering the fixes/changes genuinely validates the work with NO chance of false results or bluff of any kind. Any finding MUST be fixed, polished, improved, and covered with additional tests before the build proceeds. Multiple strong layers of checks.”): after all batch fixes/changes/implementations are done and BEFORE the pre-build test sweep + the main (artifact) build, the agent MUST dispatch dedicated code-review agent(s) (subagent-driven per §11.4.70/.20) that (1) analyze all batch work, (2) all existing data/facts/captured-evidence, (3) the existing codebase + blast radius, (4) current git history, (5) determine quality + safety + will-it-REALLY-work, (6) validate+verify every covering test genuinely exercises the work-under-test and catches its negation with ZERO false-resul bluff possibility; any finding (defect / error-prone change / safety risk / will-not-work / bluff-capable test / coverage gap) MUST be fixed+polished+improved+test-covered (four-layer §11.4.4(b), TDD-RED-first §11.4.43/.115) BEFORE the build proceeds; review iterates until no blocking findings; one of MULTIPLE STRONG LAYERS complementing (never replacing) the pre-build sweep + §11.4.92 multi-pass (author-sid + §11.4.108 + §11.4.110. Composes §11.4/.1/.4/.6/.40/.43/.50/.70/.20/.92/.102/.107/.108/.110. Propagation gate CM-COVENANT-114-125-PROPAGATION + recommended gate CM-CODE-REVIEW-GATE-BEFORE-BUILD (gate-code = separate work item). Classification: universal (§11.4.17). CLAUDE.md + AGENTS.md + QWEN.md mirrors updated in lockstep. Prior round: Added §11.4.124 Dead/unwired-code investigate-before-remove mandate (User mandat 2026-06-04, verbatim “Before removing any seemingly-dead (zero-importer / unwired) codebase, we MUST investigate via git history where/how it was originally used and ho it became dead. Removal is permitted ONLY when we have captured PROOF it is genuinely no longer needed — and that removal MUST be its own separate commit wit a proper descriptive message. If there is no such proof, the code MUST be investigated for where/how it should be wired in properly, and any missing or unwired tests MUST b added. We MUST ALWAYS be extra careful with any codebase removal.”): “zero importers ⇒ dead ⇒ delete” is a guess (§11.4.6), never a finding — before removing ANY seemingly-dead element (zero-importer / never-called / unwired function/type/file module/package/asset/config/build-target) the agent MUST FIRST investigate via git history (git log --follow, -S/-G pickaxe, blame on the deleted call-site) and capture as FACT where/how it was wired in, when/how it became dead, and whether “no references” is real or a hidden reference the static tool cannot see (reflection/DI/build-tags/codegen/plugin/FFI/config); removal is permitted ONLY with captured PROOF it i genuinely unneeded AND MUST be its own separate descriptive commit citing the git-history evidence (+ §11.4.122 operator-confirmation for end-user capabilities; tracked v

Field	Value
	§11.4.90 Obsolete); with NO such proof the element MUST NOT be removed — instead wire it in properly (restore a mistakenly-deleted call-site per §11.4.114; finish never-completed wiring) and add any missing/unwired tests (§11.4.27/.43/.115); ALWAYS be extra careful — when uncertain, default to NOT removing per §11.4.6/.101/.122. Composes §11.4.6/.8/.84/.90/.92/.101/.114/.122/.27/.43/.115. Propagation gate CM-COVENANT-114-124-PROPAGATION + recommended gate CM-DEAD-CODE-INVESTIGATE-BEFORE-REMOVE (gate-code = separate work item). Classification: universal (§11.4.17). CLAUDE.md + AGENTS.md + QWEN.md mirrors updated in lockstep. Prior round: Added §11.4.123 Rock-solid-proof-or-deep-research mandate (User mandate 2026-06-03, verbatim “Every single reported issue MUST BE fully and 100% validated with rock solid proofs! Nothing can be considered fixed or completed without hard evidence! No false results or bluff(s) of any kind is allowed! If we are not sure on how to achieve full testing, validation and verification of something we MUST ALWAYS perform deep web research ... which produces rock-solid evidence(s) and leaves no space for false results or any kind of bluff!”); every reported issue / fix / claimed completion MUST be 100% validated with rock-solid CAPTURED proof (§11.4.5/.69/.107) before it is marked fixed/implemented/completed (§11.4.33); metadata-only / config-only / absence-of-error / grep-without-runtime PASS forbidden (§11.4/.1); when UNSURE how to validate (no evidence-producing method, or only metadata/config evidence) the agent MUST ALWAYS first perform deep web research per §11.4.8+§11.4.99 to discover or build an evidence-producing validation method — declaring “untestable” or accepting a metadata-only PASS WITHOUT exhausting that research is itself a §11.4.123 violation; only after the research genuinely fails may the item be PENDING_FORENSICS:/Operator-blocked (§11.4.21)/structurally-impossible won’t-fix (§11.4.112), with the cited research as earned evidence. Forensic case study (FACT): in the 1.1.8-dev remediation the validation method for FLAG_SECURE secondary-display capture + non-introspectable streaming-app UIs was unclear; deep web research (docs/research/testing_frameworks_20260603/) yielded the CV/OCR/liveness/sink-probe oracle stack (§11.4.107/.112/.117), making rock-solid evidence possible. Composes §11.4.5/.6/.8/.52/.69/.99/.107/.118/.21/.112. Propagation gate CM-COVENANT-114-123-PROPAGATION + recommended gate CM-ROCK-SOLID-PROOF-OR-RESEARCH (gate-code = separate work item). Classification: universal (§11.4.17). CLAUDE.md + AGENTS.md + QWEN.md mirrors updated in lockstep. Prior round: Added §11.4.122 No-silent-removal-of-existing-components-without-operator-confirmation mandate (User mandate 2026-06-03, verbatim “Never ever remove any application, system component or service from already existing codebase / System without interactively asked question to us! THIS IS MANDATORY RULE / CONSTRAINT!”); forensic case study — F2 Apple-TV-class app + F4 Huawei HMS component removed in the 1.1.8-dev burn-down WITHOUT asking, operator reversed both); no application/component/service/package/feature/driver/module/library/prebuilt — any already-existing end-user capability — may be removed from the existing codebase/System without FIRST interactively asking the operator (§11.4.66 mechanism never silent, never autonomous) and receiving an EXPLICIT keep-or-remove decision; silent removal is a release blocker regardless of rationale; tracked DROP path is ask→operator-approves→mark Obsolete (→ Fixed.md) reason feature-removed+operator-approval citation (§11.4.90)→then remove. Composes §11.4.66/.101/.90/.112/.6/.40/.42. Propagation gate CM-COVENANT-114-122-PROPAGATION + recommended gate CM-NO-SILENT-COMPONENT-REMOVAL (gate-code = separate work item). Classification: universal (§11.4.17). CLAUDE.md + AGENTS.md + QWEN.md mirrors updated in lockstep. Prior round: Added §11.4.114–§11.4.121 (eight new universal anchors from the 1.1.8-dev live-defect remediation, 2026-06-03): §11.4.114 last-known-good-tag regression isolation (diff the broken state against the known-good tag FIRST — the tagged-good version is the regression oracle; surgical

Field	Value
	forward-fix over wholesale revert); §11.4.115 RED-baseline-on-the-broken-artifact + polarity-switch (one test reproduces the defect on the pre-fix artifact then flips RED_MODE=0 to GREEN-guard — refines §11.4.43’s RED step); §11.4.116 real-time conductor[?]autonomous-test-framework sync channel (append-only JSONL event stream + atomically-rewritten status snapshot the orchestrator tails live; every verdict event carries its evidence path, project-neutral per §11.4.28); §11.4.117 CV/OCR pixel-oracle fallback for non-introspectable UIs (drive + assert via pixels when the accessibility tree is blank — TV-Compose / leanback / canvas / games; refines §11.4.48/.52/.107); §11.4.118 discovery-pressure to confirm known-issue-set completeness (fixing the reported set is necessary-not-sufficient; provable enumerated discovery coverage, not “found nothing else”); §11.4.119 single-resource-owner partitioning for parallel hardware testing (exactly ONE stream owns each device’s exclusive resource, others read-only — refines §11.4.58/.103 for hardware contention); §11.4.120 fix-breaks-its-own-gate reconciliation (the gate FAIL is the correct signal the fix worked — rewrite the gate to assert the NEW mechanism + update its mutation; never fake-pass, never revert the fix) §11.4.121 no-commit-while-build-writes-tracked-artifacts (defer commit to build completion so partial/stale artifacts never land in VCS — build-output analogue of §11.4.84). All eight Classification: universal (§11.4.17). CLAUDE.md + AGENTS.md + QWEN.md mirrors updated in lockstep. Prior round: §11.4.108 Four-layer fix-verification + runtime-signature-as-definition-of-done mandate (systematic-debugging Phase 4.5, 2026-06-03). A fix crosses FOUR distinct layers — SOURCE (committed) — ARTIFACT (bytes landed in the produced artifact) → RUNTIME-ON-CLEAN-TARGET (active on a clean/fresh deployment) → USER-VISIBLE (works for the end user) — and “fixed” at one layer does NOT imply fixed at the next; source-green is the cheapest least conclusive signal. Every fix declares ONE machine-checkable runtime signature verified on a clean target; this per-fix runtime-signature registry is the SINGLE SOURCE OF TRUTH for “fixed”, REPLACING “a gate greps the source”. Gates span all four layers (artifact gate asserts BYTES landed, not merely that the source still contains the change). Eliminate the stale-deployment/shadow layer by construction — deployment MUST yield a clean state (wipe the mutable overlay) OR pre-validation MUST prove running-artifact == built-artifact; validation against possibly-stale state is INVALID. Meta-rule: ≥3 “fixed-but-not-working” discoveries in one cycle = ONE architectural VERIFICATION flaw (not independent bugs) → STOP patching symptoms, fix the pipeline, re-certify every item on a clean target. A batch is validated only after comprehensive per-item runtime-signature verification on a clean baseline. Composes §11.4.1/.2/.4/.5/.6/.27/.40/.46/.50/.52/.69/.102. Propagation gate CM-COVENANT-114-108-PROPAGATION + recommended per-fix gate CM-RUNTIME-SIGNATURE-REGISTRY (gate-code = separate work item). Classification: universal. CLAUDE.md + AGENTS.md + QWEN.md mirrors updated in lockstep. Prior round: §11.4.107 Anti-bluff AV/test-validation techniques mandate (User-driven research, 2026-06-02). A single captured frame is NOT proof of playback — prove live advancing frames via a freeze-detection oracle (near-duplicate / perceptual-hash adjacent-frame distance, NOT byte-identical) + an independent frame-advance counter from the platform compositor/decoder telemetry; loading/buffering is a distinct state (wait for genuine playback-start, never false-FAIL a still-loading stream nor false-PASS a spinner); not-stale-from-previous cross-check; no-flash-on-wrong-output high-frequency sampling during routing transitions; feed/UI-driven (not deep-link) test method; metamorphic relations for the oracle problem on un-owned streamed content; full-reference quality metrics (SSIM/VMAF/ΔE2000) vs golden source for owned content; mutation-test even analyzer with a golden-good + golden-bad fixture pair so the analyzer itself cannot bluff per-channel audio loudness (EBU R128) + XRUN census; confidence-floored ROI OCF thresholds calibrated on the project’s own fixtures (§11.4.6). Composes §11.4.5/.6/.50/.68/.69/.85. Classification: universal. CLAUDE.md + AGENTS.md +

Field	Value
	QWEN.md mirrors updated in lockstep. Prior round: §11.4.103 Continuous parallel-stream working routine (User mandate 2026-05-29) — parallel-stream floor raised ≥ 2 — ≥ 3 + auto-backfill rule (the moment a stream is fully done, a new one immediately takes its place; count never drops below 3 while actionable items remain; idle below 3 only when no remaining items OR all externally blocked). Original: the proven multi-stream operating pattern promoted to the project's standing default working routine: main work stream stays FREE (ALL commit AND push detached, never blocking); ≥ 3 parallel subagent-driven background streams whenever non-contending actionable work exists; most-critical + most-visible first with audio always top (§11.4.72); §11.4.96 safe-during build scope only (NO second gradle/m -j/containerised build while the 42 GB AOSP rebuild runs, §12.8); heavy anti-bluff on every closure (§11.4.5/.6/.50/.69/.102); idle ONLY when genuinely externally blocked (§11.4.94/.97). Composes §11.4.58/.70/.72/.87/.88/.89/.94/.96/.97/.101/.102. Propagation gate CM-COVENANT-114-103-PROPAGATION (gate-code = separate work item). Classification: universal. CLAUDE.md + AGENTS.md + QWEN.md mirrors updated in lockstep. Prior round: §11.4.102 Mandatory systematic-debugging activation + always-loaded skill-discovery + plugin-dependency availability (User mandate 2026-05-29). On ANY spotted issue/bug/test-failure/misalignment/inconsistency the agent MUST launch superpowers:systematic-debugging (or platform-equivalent) BEFORE proposing any fix — Iron Law “NO FIXES WITHOUT ROOT CAUSE INVESTIGATION FIRST”, full 4-phase arc; superpowers:using-superpowers skill-discovery ALWAYS loaded + consulted before any task (any $\geq 1\%$-applicable skill MUST be invoked); every mandated skill plugin/dependency MUST be installed + loadable (missing = release-blocker). Composes §11.4.4/.6/.7/.8/.43/.70/.82(A)/.92. Propagation gate CM-COVENANT-114-102-PROPAGATION (gate-code = separate work item). Classification: universal. CLAUDE.md + AGENTS.md + QWEN.md mirrors updated in lockstep. Prior round: §11.4.101 Autonomous-decision-over-blocking mandate (User mandate 2026-05-28).
Issues	none
Issues summary	ToC is currently stale (missing entries for §11.4.71/.72/.73/.74/.75/.76) — separate clear up commit needed per §11.4.61.
Fixed	§11.4.108 added; §11.4.73, §11.4.74, §11.4.76; QWEN.md created
Fixed summary	Operator-mandated additions 2026-05-20 propagated into Constitution.md (this commit CLAUDE.md + AGENTS.md + new QWEN.md mirrors land in the same commit. §11.4.76 renumber-from-§11.4.75 documented inline (collision with concurrent Mechanical Enforcement landing).
Continuation	(1) ToC regeneration commit (cover §11.4.71..§11.4.76). (2) Submodule-catalogue inventory expansion in CLAUDE.md / AGENTS.md / QWEN.md per the 2026-05-20 follow-up user mandate — per-Submodule purpose lines for every repo in vasic-digital + HelixDevelopment.

Table of contents

- 1. Test coverage is mandatory for every change
 - 1.1 False-positive immunity is an invariant
- 2. Commit and push mechanics — single entryptpoint, locked
 - 2.1 Multi-upstream push is the norm
- 3. Submodule changes propagate through submodule commits first
- 4. Every tag on the main repo MUST be mirrored on every owned submodule

- 5. Changelog discipline and multi-format export
- 6. Documentation up to the nano-details
- 7. Making false-success results literally impossible
 - §7.1 NO BLUFF — positive-evidence-only validation
- §8. Bleeding-edge ultra-perfection quality bar
- §9. Absolute codebase and data safety — zero risk, zero loss
 - §9.1 Mandatory safety protocol for destructive operations
 - §9.2 Force-push requires explicit user authorization every time
 - §9.3 Hardlinked backup is the standard — there is no excuse
 - §9.4 Commit-message audit trail for history rewrites
- §10. Enforcement
- §11. End-user quality covenant
 - §11.4 End-user quality guarantee — forensic anchor (User mandate, 2026-04-28)
 - §11.4.1 — FAIL-bluffs are equally forbidden
 - §11.4.2 — Recorded-evidence requirement
 - §11.4.3 — Per-environment-topology test dispatch
 - §11.4.4 — Test-interrupt-on-discovery + retest-from-clean-baseline
 - §11.4.5 — Captured-evidence quality analysis
 - §11.4.6 — No-guessing mandate
 - §11.4.7 — Demotion-evidence rule
 - §11.4.8 — Deep-web-research-before-implementation
 - §11.4.9 — Batch-source-fixes-before-rebuild
 - §11.4.10 — Credentials-handling mandate
 - §11.4.10.A — Pre-store credential leak audit (User mandate, 2026-05-17)
 - §11.4.11 — File-layout discipline
 - §11.4.12 — Auto-generated docs sync mandate
 - §11.4.13 — Out-of-band sink-side captured-evidence
 - §11.4.14 — Test playback cleanup mandate
 - §11.4.15 — Item-status tracking mandate
 - §11.4.16 — Item-type tracking mandate
 - §11.4.17 — Universal-vs-project classification of new rules (User mandate, 2026-05-14)
 - §11.4.20 — Subagent-driven-by-default mandate (User mandate, 2026-05-14)
 - §11.4.18 — Script documentation mandate (User mandate, 2026-05-14)
 - §11.4.19 — Fixed-document column-alignment mandate (User mandate, 2026-05-14)
 - §11.4.21 — Operator-blocked status + self-resolution exhaustion mandate (User mandate, 2026-05-14)
 - §11.4.22 — Document-sync commit discipline (User mandate, 2026-05-14)
 - §11.4.23 — Visual-cue & grouping mandate for Issues docs (User mandate, 2026-05-14)
 - §11.4.24 — Build-resource stats tracking mandate (User mandate, 2026-05-14)
 - §11.4.25 — Full-Automation-Coverage Mandate (User mandate, 2026-05-15)
 - §11.4.26 — Constitution-Submodule Update Workflow Mandate (User mandate, 2026-05-15)
 - §11.4.31 — Submodule-Dependency-Manifest Mandate (User mandate, 2026-05-15)
 - §11.4.32 — Post-Constitution-Pull Validation Mandate (User mandate, 2026-05-15)

- §11.4.30 — .gitignore + No-Versioned-Build-Artifacts Mandate (User mandate, 2026-05-15)
- §11.4.29 — Lowercase-Snake Case-Naming Mandate (User mandate, 2026-05-15)
- §11.4.28 — Submodules-As-Equal-Codebase + Decoupling + Dependency-Layout Mandate (User mandate, 2026-05-15)
- §11.4.27 — No-Fakes-Beyond-Unit-Tests + 100%-Test-Type-Coverage Mandate (User mandate, 2026-05-15)
- §11.4.33 — Type-aware closure-status vocabulary (User mandate, 2026-05-15)
- §11.4.34 — Reopened-source attribution mandate (User mandate, 2026-05-15)
- §11.4.35 — Canonical-root inheritance clarity (User mandate, 2026-05-15)
- §11.4.36 — Mandatory install_upstreams on clone/add Mandate (User mandate, 2026-05-15)
- §11.4.37 — Fetch-before-edit mandate (User mandate, 2026-05-15)
- §11.4.38 — Installable-Asset Evidence Mandate (User mandate, 2026-05-17)
- §11.4.39 — Per-Feature On-Device End-User Validation Mandate (iter-76, 2026-05-17)
- §11.4.40 — Full-suite retest before release tag mandate (User mandate, 2026-05-17)
- §11.4.41 — Pre-Force-Push Merge-First Mandate (User mandate, 2026-05-17)
- §11.4.42 — Iteration-discipline mandate (User mandate, 2026-05-18)
- §11.4.43 — TDD-Fix-Discipline Mandate (User mandate, 2026-05-18)
- §11.4.44 — Document Revision Header Mandate (User mandate, 2026-05-18)
- §11.4.45 — Integration-Status-Doc Maintenance Mandate (User mandate, 2026-05-18)
- §11.4.46 — Validate-recent-work-before-post-flash-tests mandate (User mandate, 2026-05-18)
- §11.4.47 — Firebase Data Review Mandate (User mandate, 2026-05-18)
- §11.4.48 — UI-Driven Video Testing Mandate (User mandate, 2026-05-18)
- §11.4.49 — Dual-Approach Testing Mandate (User mandate, 2026-05-18)
- §11.4.50 — Deterministic Consistency Mandate (User mandate, 2026-05-18)
- §11.4.51 — Live-ADB-First Maximization Mandate (User mandate, 2026-05-18)
- §11.4.52 — Autonomous-Validation Mandate (User mandate, 2026-05-18)
- §11.4.53 — Fixed Summary parity mandate (User mandate, 2026-05-18)
- §11.4.54 — ATM-NNN ticket identifier mandate (User mandate, 2026-05-19)
- §11.4.55 — Reopens-history tracking + per-item Reopens.md doc (User mandate, 2026-05-19)
- §11.4.56 — Status_Summary parity + two-audience format (User mandate, 2026-05-19)
- §11.4.57 — README.md doc-link section + revision metadata (User mandate, 2026-05-19)
- §11.4.58 — Parallel-development methodology (User mandate, 2026-05-19)
- §11.4.59 — README always-sync mandate (User mandate, 2026-05-19)
- §11.4.60 — Documentation always-sync composite covenant (User mandate, 2026-05-19)

- [§11.4.61 — Mandatory Markdown metadata table + structured-doc ToC \(User mandate, 2026-05-19\)](#)
- [§11.4.63 — Workable-items procedure docs as single source of truth \(User mandate, 2026-05-19\)](#)
- [§11.4.65 — Universal Markdown export mandate \(User mandate, 2026-05-19\)](#)
- [§11.4.66 — Blocker-resolution interactive-clarification mandate \(User mandate, 2026-05-19\)](#)
- [§12. Host-session safety — directly OR indirectly signing the user out is FORBIDDEN](#)
 - [§12.1 Forbidden operations — directly OR indirectly](#)
 - [§12.2 Required safeguards](#)
 - [§12.3 Container hygiene](#)
 - [§12.6 Memory-Budget Ceiling — 60% MAXIMUM](#)
 - [§12.10 Continuation document — sacred invariant](#)
- [Appendix A — Mutation testing — academic and industrial foundations](#)
- [Appendix B — Recursive inheritance & path-independence](#)
- [Appendix C — Multi-upstream push](#)

This document defines mandatory, non-negotiable rules for every project that includes this constitution submodule. Every AI agent and every human contributor MUST comply. Violations are blockers.

Project-specific constitutions extend this document. When a project Constitution (e.g. docs/guides/PROJECT_CONSTITUTION.md) and this universal Constitution differ, the project Constitution may extend or tighten — but NEVER weaken — the universal rules below.

Last revision: 2026-05-14 (v1.0 — initial extraction from ATMOSphere Android 15 1.1.5-dev)

This Constitution is the universal source-of-truth for the engineering practices shared across all Helix projects (and any project that opts in by adding this submodule). It is intentionally agnostic about specific hardware, operating systems, languages, vendors, and product domains. Anything in this document that mentions a specific product, hardware SKU, library version, or vendor MUST be moved to the consuming project's own Constitution.

1. Test coverage is mandatory for every change

Every code, configuration, documentation, or infrastructure change MUST ship with tests that prove:

Invariant	Mechanism
The change is present in source	pre-build / pre-merge gate that scans the source tree for the expected pattern
The change survives compilation / packaging	post-build gate that inspects the produced artifact (binary, archive, image) and verifies the change actually shipped

Invariant	Mechanism
The change behaves correctly at runtime	runtime / integration / on-device test that exercises the user-visible behaviour
The gate itself is not bluffing	meta-test (mutation) entry that mutates the asserted condition and proves the gate catches the break

A change without all four layers of coverage is not ready to merge.

1.1 False-positive immunity is an invariant

A test that always returns PASS because its regex never matches, its input path is wrong, its assertion target does not exist, or its comparison is tautological is **worse than no test**. Every new gate **MUST** be paired with a mutation entry in the project's meta-test harness that:

1. Temporarily breaks the assertion (sed out a line, rename a symbol, etc.).
2. Re-runs the gate.
3. Asserts the gate now reports FAIL.
4. Restores the original.

If the mutation round does not turn PASS → FAIL, the gate is a sham and must be rewritten.

This is the **single most important rule in this Constitution**. Every subsequent §11.4.x clause is downstream of it.

2. Commit and push mechanics — single entriypoint, locked

All commit and push work uses the project's official multi-remote commit wrapper (e.g. `scripts/commit_all.sh`) as the single entriypoint. Direct `git commit` / `git push` / `git add` on the main repo is prohibited in normal workflow. Exception: tag creation + tag-ref push is done via explicit `git tag -a` and `git push <remote> <tag>`.

The commit and push wrappers **MUST** hold an advisory `flock` so two invocations cannot race against each other. Lock files self-clean on process exit via `trap`.

When a contributor sees the “another commit/push wrapper is already running” error, they **MUST NOT** `rm -f` the lock unless they have just killed the owning process.

2.1 Multi-upstream push is the norm

Every project hosted on multiple Git providers (GitHub primary, GitLab / GitFlic / GitVerse / Bitbucket / Gitea mirrors) **MUST** push every commit to ALL configured upstream remotes. A commit that lives on only one remote is a future operational risk: when that one provider has an outage, the project is unreachable.

The constitution submodule itself ships an `install_upstreams.sh` helper (and an `Upstreams/` directory with one declaration script per remote). Running it from the submodule root configures all the remotes locally. Consuming projects **SHOULD** do the same for their own multi-remote topology.

3. Submodule changes propagate through submodule commits first

When a change lands in any submodule of the project:

1. **Commit inside the submodule first** (the submodule's own `git add + git commit`), since the parent project's commit wrapper does not commit submodule source — it only captures the updated submodule pointer.
2. **Push the submodule commit to all** remotes of that submodule.
3. **Then** run the parent project's commit wrapper to capture the updated submodule pointer in main and push main.

Skipping step 1 produces parent commits / tags that point at old submodule HEADs without the actual source.

4. Every tag on the main repo **MUST** be mirrored on every owned submodule

When a version tag is created on the main repo at a specific commit, the same tag **MUST** be created on **every owned submodule** at that submodule's currently-pointed-to HEAD, and pushed to **every remote** of every owned repo. Third-party submodules (libraries not under the project's control) are excluded. The consuming project's Constitution declares its owned-submodule set explicitly.

5. Changelog discipline and multi-format export

Every tagged release **MUST** ship:

1. A new `docs/changelogs/<tag>.md` entry describing user-visible changes, developer-visible changes, and known caveats.
2. Exports of that entry to `docs/changelogs/<tag>.{html,json,txt}` so external tooling (release portals, ticketing systems, package indices) can consume without needing a Markdown parser.

3. An update to a cumulative docs/change logs/CHANGELOG.md (reverse chronological).

The project SHOULD provide a script (scripts/testing/export_changelog.sh <tag> or equivalent) that wraps pandoc + jq + plain-text Markdown strip.

6. Documentation up to the nano-details

Every new feature, fix, or infrastructure change MUST update:

1. The project's CLAUDE.md / AGENTS.md Applied Fixes table (one row, one commit, one release tag).
2. docs/guides/ for any user-visible or developer-reachable behaviour change.
3. Architecture diagrams, flowcharts, and any reference docs that touch the changed subsystem.
4. The per-version changelog (§5).

Documentation drift after a fix is itself a Constitution violation: undocumented behaviour change is the same defect class as un-tested behaviour change.

7. Making false-success results literally impossible

All validation output MUST satisfy the following invariants:

1. Every gate reports concrete PASS / FAIL / SKIP with explicit reason text. A gate that just says “ok” is non-compliant.
2. Every SKIP reason MUST be mechanically distinguishable from a PASS. Summary counters track them separately.
3. FAILs are counted towards exit status. A script that hides FAILs in logs and returns 0 is a bug.
4. Meta-tests (§1.1) catch any gate that always PASSES regardless of condition.
5. Runtime results MUST be cross-referenced against host-side evidence (capture rig, screen recording, log analyzer, or equivalent) when the system's own introspection cannot directly verify the behaviour.

§7.1 NO BLUFF — positive-evidence-only validation

Every runtime test MUST satisfy ALL FIVE of the following non-negotiable constraints. A test that violates any one of them is a **bluff** and is forbidden from being committed to the test suite:

1. **Real ACTION**: the test MUST invoke at least one user-visible action via a project anti-bluff helper (e.g. `ab_send_action()` in shell-based suites). A test that records a PASS without ever calling such a helper is automatically failed by the summary function.
2. **State DELTA**: every PASS that claims to verify a feature MUST capture state BEFORE the action and AFTER, and assert the delta matches expectation. `before == after` on success is a bluff.

3. **POSITIVE EVIDENCE:** the PASS condition **MUST** be POSITIVE evidence (a value match, a state delta, a captured frame analysis, a captured audio frame analysis), **NEVER** absence-of-error. If a probe returns nothing, that's a FAIL not a SKIP.
4. **UNIQUE EVIDENCE TOKEN:** tests that interact with mutable framework state **MUST** embed a per-run UUID into the action **AND** look for that token in the resulting state. Cached results predating the token cannot match — this defeats the stale-cache false-pass.
5. **Audio / video features REQUIRE captured evidence:** features that produce audio output **MUST** be validated by capturing the actual output via a capture pipeline (loopback, hardware capture rig, or equivalent) — **NOT** by metadata-only checks. Features that produce video output **MUST** be validated via frame capture + OCR or pixel-difference analysis.

Each consuming project **SHOULD** ship a shared anti-bluff helper library (typically `tests/lib/anti_bluff.sh` or equivalent) that codifies these rules; every new runtime test **MUST** source it and use its assertion helpers.

§8. Bleeding-edge ultra-perfection quality bar

The acceptance bar for shipping a project change is **not** “tests pass and source compiles” — that is the floor. The acceptance bar is:

The user, on a freshly-deployed system, sees the expected effect when they perform the documented action — confirmed end-to-end, with captured evidence, on the actual target environment.

The following are **forbidden**:

1. **“Source rebrand without shipped artifact rebuild”:** every source-side rename or rebranding **MUST** flow through to a shipped artifact that the runtime actually loads. Pre-build source-string gates are insufficient by themselves; shipped-artifact gates must accompany them.
 2. **“Tests pass but feature broken”:** this is the failure mode §7.1 exists to eradicate. Any test that **PASSES** despite the feature being non-functional is a critical defect; its mutation pair must **FAIL** on the breakage.
 3. **“Configuration-only tests”:** a test that only verifies a config file exists / has expected contents is downgraded to a pre-build gate. Runtime tests **MUST** exercise the actual user-visible behaviour and capture the result.
 4. **“It works on my workstation”:** every fix is validated on every supported target before any release. No exceptions.
 5. **“Will fix in next cycle”:** deferrals are documented in the changelog with concrete blockers + implementation plans, not just handwaved.
-

§9. Absolute codebase and data safety — zero risk, zero loss

Every destructive operation on the repository (history rewrite, force-push, branch deletion, bulk file removal, submodule de-init, pack repack that drops objects) is a **safety-critical** operation. Data loss from a wrong force-push is irreversible once the remote garbage-collects dangling objects. This section is non-negotiable.

§9.1 Mandatory safety protocol for destructive operations

Every destructive operation **MUST** execute, in strict order:

1. **Full backup before touching anything.** Hardlinked mirror of `.git` to a sibling backup directory:

```
cp -al .git <backup>/<repo>.git.mirror
```

This is near-instant on the same filesystem and uses zero additional disk. If `.git` is small, also create `git bundle create <backup>/main.bundle --all`.

2. **Record critical metadata into the backup dir:**

- `git show-ref > refs.txt`
- `git tag > tags.txt`
- `git submodule status > submodules.txt`
- `git rev-parse HEAD > head.txt`
- `git rev-parse HEAD^{tree} > head_tree.txt`
- `git ls-tree -r HEAD --full-tree | sha256sum > head_tree_content.sha256`
- `git ls-tree -r HEAD --full-tree > head_tree_listing.txt`

3. **Identify the target state:** the expected HEAD commit, tree hash, and tree content hash after the operation completes. The operation **MUST** preserve these unless it is explicitly changing HEAD content.

4. **Run the operation** (filter-repo, rebase, etc.). NEVER use `--no-verify / bypass-hooks` flags. NEVER auto-force on failure.

5. **Post-operation verification gate** (all **MUST** pass):

- HEAD tree hash matches pre-op target (content identical for non-content-changing rewrites)
- HEAD tree content sha256 matches pre-op sha256
- All tags from `tags.txt` are preserved and resolve to valid commits
- All submodule pointers from `submodules.txt` match current state
- Project's pre-build / pre-merge gates pass with zero FAIL and zero unclassified WARN

6. **Only if every check in §9.1.5 passes** may the operation be considered successful. Otherwise: restore from backup immediately
(`rm -rf .git && mv <backup>/<repo>.git.mirror .git`).

7. **Then and only then** may a force-push proceed — and still not without explicit user authorization per §9.2.

§9.2 Force-push requires explicit user authorization every time

Force-pushing (including `push --force`, `push --force-with-lease`, or any divergence-recovery code path inside an automated wrapper) overwrites the remote's view of history. On any shared branch this is irreversible.

Rules: - Automated push wrappers **MUST NOT** auto-force-push as a fallback to any kind of failure (size rejection, divergence, server error, anything). The historical divergence-recovery code that silently force-resets the remote is forbidden and must be removed or gated behind an explicit flag. - A human must affirmatively say “force-push” (or equivalent) in the session to authorize each force-push. - Force-push to main (and any equivalent primary branch on owned submodules) is only authorized after the §9.1.5 gate has passed. - Every force-push event is recorded in `docs/changelogs/<tag>.md` under a “Force-push audit” section (target refs, backup location, validation gate output).

§9.3 Hardlinked backup is the standard — there is no excuse

Because hardlinked `.git` copies are near-instant and use zero additional disk (same-filesystem inodes), the cost of making a backup is trivial. Any destructive operation without a fresh hardlinked backup is a Constitution violation regardless of how small the repo is or how confident the operator.

§9.4 Commit-message audit trail for history rewrites

After any `git filter-repo` run (or equivalent), commit a record under `docs/changelogs/` (or `docs/history-rewrites/`) containing: - What was stripped and why - Pre-op `.git` size vs post-op `.git` size - Pre-op HEAD commit vs post-op HEAD commit - Pre-op HEAD tree hash vs post-op HEAD tree hash (MUST match for non-content-changing rewrites) - Backup location

§10. Enforcement

Any commit, tag, or release that violates §1–§9 is non-compliant. The fix is to amend/revert and re-land with compliance. No exceptions for speed pressure. Data-safety violations (§9) and host-session-safety violations (§12) are treated as catastrophic and block the entire release cycle until fully remediated.

§11. End-user quality covenant

§11.4 End-user quality guarantee — forensic anchor (User mandate, 2026-04-28)

The reason §7.1 + §8 exist — captured verbatim from the project owner so that future engineers and AI agents understand the historical failure mode the covenant exists to eradicate:

“We had been in position that all tests do execute with success and all Challenges as well, but in reality the most of the features does not work and can’t be used! This **MUST NOT** be the case and execution of tests and Challenges **MUST** guarantee the quality, the completion and full usability by end users of the product!”

This is the canonical motivation for every gate, every mutation pair, every meta-test, every anti-bluff helper, every captured- evidence requirement, every multi-environment validation rule. The covenant exists because the project has historically shipped broken features behind PASS-ing tests, and that outcome is no longer tolerated.

Operative rule: the bar for shipping is not “tests pass” but “**users can use the feature.**” Every PASS in this codebase **MUST** carry positive evidence captured during execution that the feature works for the end user. Metadata-only PASS, configuration-only PASS, “absence-of-error” PASS, and grep-based PASS without runtime evidence are all critical defects regardless of how green the summary line looks.

Propagation requirement: every `CLAUDE.md` and every `AGENTS.md` in the project tree (parent + every submodule recursively) **MUST** contain a clearly-titled “**MANDATORY ANTI-BLUFF COVENANT — END-USER QUALITY GUARANTEE**” section that quotes this user mandate verbatim and points back to this §11.4 as the canonical authority. The consuming project enforces this via a pre-build gate (typically `CM-COVENANT-PROPAGATION`).

§11.4.1 — FAIL-bluffs are equally forbidden

The covenant must be enforced symmetrically. The historical failure mode the covenant eradicates is the **PASS-bluff**: a green test result on a feature that doesn’t work for users. But the inverse failure mode is just as toxic — the **FAIL-bluff**: a red test result caused by a script bug (shell crash, missing argument, regex error, malformed assertion) rather than by an actual product defect.

Examples of FAIL-bluffs: - A test using `set -eu` calls `log_pass` without arguments → the helper references `$1` → script crashes with “1: parameter not set” → orchestrator records FAIL → engineer assumes product bug. - A test pre-condition check uses an unbounded `sed` range → a comment line in the source matches the range terminator → the intended mutation never fires → audit reports PASS — bluff.

A FAIL produced by a test bug is no more useful than a PASS produced by a test bug — both mislead investigation, waste cycles, and ultimately let real defects ship undetected.

Operative rule (extended): every test MUST fail ONLY for genuine product defects. A test that crashes for a script-internal reason — undefined variable under set -u, unhandled error, malformed pipe, regex syntax error, division by zero, etc. — is a critical defect in the test suite itself and MUST be fixed at the source layer (helper library, shared lib, test source) so the failure mode disappears project-wide, not patched in individual call sites.

§11.4.2 — Recorded-evidence requirement

A test that emits PASS without **captured visual or audio evidence of the user-visible feature actually working on the screen the user would see** is a §11.4 PASS-bluff. Closing this gap requires a project-wide recording + analyzer infrastructure consisting of (at minimum):

- **Recording wrapper** that captures the user-visible output to time-synchronized media files. Multi-display systems record every output stream in parallel; sync metadata (host wall-clock at spawn) is written alongside the recording.
- **Action-timeline emitter:** every test emits structured timeline events (JSONL or equivalent) at every user-visible transition.
- **Frame / audio analyzer:** scans the recordings, extracts samples at a configurable interval, feeds them to recognizers (OCR for text overlays, speech-to-text for audio, pixel-diff for frame health), correlates against the timeline, and emits findings.

Closure criterion: every PASS for a user-visible feature MUST be cross-checked by the analyzer against the recording + action timeline. A PASS that lacks at least one matched timeline event in the analyzer findings is treated as a §11.4 PASS-bluff.

§11.4.3 — Per-environment-topology test dispatch

Tests that depend on environment topology (presence/absence of secondary displays, microphones, network interfaces, peripheral devices, cloud services, etc.) MUST detect topology at test entry and dispatch the topology-appropriate variant. A test that runs the WRONG variant for the actual topology and PASSes is a §11.4 PASS-bluff: it claims a feature works while never having exercised it on the configuration the user actually has.

A topology-touching test that does NOT have a dispatcher AND a per-topology variant is a §11.4 violation regardless of how it “passes” on any specific configuration. Single-test-multi-topology scripts MUST gate every topology-dependent assertion behind an explicit topology check that emits SKIP-with-reason if the required topology is absent — never PASS-by-default.

§11.4.4 — Test-interrupt-on-discovery + retest-from-clean-baseline

A test cycle that *continues running past a freshly discovered defect* is itself a §11.4 PASS-bluff: it produces “all green” summaries that the operator might tag as a release while the codebase under test is known-broken at the moment those greens were recorded.

The mandatory protocol — non-negotiable:

1. **Interrupt immediately.** The moment a defect is re-discovered, re-produced, or newly identified during a test cycle, the cycle **MUST** stop. No “let it finish for the data” — the data is contaminated.
2. **Systematic debugging before any fix.** Identify root cause at source layer (not symptom), blast radius across related tests/features/subsystems, and regression-protection seam. Symptom patching is forbidden.
3. **Fix at root cause** per §11.4.1: source-layer fix, not symptom patch in the calling site. The fix **MUST** be safe (no new failure modes), minimal (no surrounding cleanup smuggled in), and non-regressive.
4. **Four-layer test coverage per fix.** Every fix lands with positive-evidence coverage in every applicable layer:
 - **Pre-build / pre-merge gate** (source-level catch)
 - **Post-build / packaging gate** (artifact-level catch)
 - **Runtime / on-device test** (end-user-behaviour catch)
 - **Meta-test paired mutation** (proves the gate is not itself a bluff)
5. **Documentation **MUST** be updated for every fix** — Issues → Fixed migration on closure, Applied Fixes table row, user-facing guides, architecture diagrams, per-version changelog.
6. **Rebuild the system.** Full build via the canonical build script. Skipping rebuild — even when the fix is host-only — invites stale-artifact contamination of the retest baseline.
7. **Re-deploy on every supported target.** A target left on the pre-fix build is a confounder.
8. **Repeat the full test suite** from the beginning on every target, sequentially per host memory cap. Partial reruns are a §11.4 violation.
9. **No-bluff certification per cycle.** Before tagging: meta-test harness returns all gates green **AND** every gate’s paired mutation FAILs.

§11.4.5 — Captured-evidence quality analysis

§11.4.2 mandates *captured* evidence; §11.4.5 mandates the **content** of that evidence be analyzed for quality, not merely for presence.

Audio quality analysis — every audio test that **PASSes** **MUST** verify: 1.

Presence — non-trivial RMS amplitude in captured output. 2. **Channel count** — analyzer reports the exact channel count the test claims (2.0 stereo, 5.1 surround, etc.). Stereo downmix in a “5.1 PASS” claim is a PASS-bluff. 3. **Sample rate + bit depth** — match the codec / pipeline under test. 4. **Glitch census** — underruns, XRUNs, dropouts above tolerance **MUST** be explicitly classified (PASS within budget, WARN above, FAIL on hard limits). Silently ignoring counts is a PASS-bluff. 5. **Coexistence-artifact census** — radio / IPC / shared-resource contention above tolerance.

Video quality analysis — every video test that **PASSes** **MUST** verify: 1.

Presence — captured recording has non-zero size **AND** decoded- frame total > 0. A 0-byte file is the canonical PASS-bluff. 2. **Routing target** — analyzer confirms video appeared on the intended display / surface. 3. **Frame health** — drop count, jitter, freeze detection, tearing. 4. **Obstruction census** — OCR scan for hostile overlays (error dialogs, sign-in dialogs, geo-restriction overlays, paywall, etc). 5. **Resolution + codec** — captured dimensions match what the test claims.

§11.4.6 — No-guessing mandate

Tests, gates, status reports, closure narratives, commit messages, and any operator-facing text **MUST NOT** use words like likely, probably, maybe, might, possibly, presumably, seems, appears to, guess, seemingly, apparently, perhaps, supposedly, conjectured, or their synonyms when describing **CAUSES** of test failures, system behaviour, or fix effectiveness.

Either:

1. **Prove the cause** with captured forensic evidence (logs, kernel ring buffer, kernel ramoops, structured snapshots, OS journals, strace, etc.) and state it as fact, OR
2. **Explicitly mark UNCONFIRMED + PENDING_FORENSICS** with a tracked-task ID for follow-up.

Every “X likely caused Y” sentence in the codebase or documentation is a §11.4.6 violation. Every “appears to be benign” without a concrete forensic trace is a §11.4.6 violation.

§11.4.7 — Demotion-evidence rule

A demotion from any FAIL classification (OPEN, POSSIBLE PRODUCT DEFECT, FAIL) to a lower-severity classification (INVESTIGATED, MITIGATED, RESOLVED, WORKING-AS-INTENDED) requires positive evidence captured under the **SAME CONDITIONS** that originally exposed the defect:

1. **Same target** — defect on Target-A needs Target-A evidence; cross-target claims need every target.
2. **Same build** — evidence captured BEFORE a relevant fix landed does not validate post-fix behaviour.
3. **Same cycle position** — a stress-soak FAIL is not refuted by an isolated-test PASS.
4. **Same load profile** — a contention-mode FAIL is not refuted by a quiescent-mode PASS.

“I cannot reproduce in isolation” is a HYPOTHESIS, not a finding. Per §11.4.6 it **MUST** be tagged **UNCONFIRMED**: until same-conditions retest produces positive evidence.

§11.4.8 — Deep-web-research-before-implementation

Before designing a non-trivial fix, before implementing a new feature, before declaring an architectural choice — perform deep web research to verify the chosen approach is informed by current state-of-the-art. The research surface includes:

1. **Official documentation** of the platform, framework, language, or service.
2. **Vendor technical guides** for the hardware / cloud / library.
3. **Open-source codebases** that already solved analogous problems.
4. **Coding tutorials + technical articles**.
5. **Issue trackers** for the relevant projects.

Operative rule. A fix that re-invents a wheel (or reproduces a known-broken pattern) when the open-source community has already solved the problem is a §11.4 violation by omission.

Documentation requirement. Every non-trivial fix’s commit message (or accompanying entry in the project’s Issues / Fixed file) MUST cite the research sources that informed it: at least one external link OR the literal “NO external solution found — original work”.

§11.4.9 — Batch-source-fixes-before-rebuild

When working through a multi-defect closure cycle, all source-side fixes that DO NOT require runtime validation to design MUST be landed BEFORE the next artifact rebuild. The anti-pattern this mandate eliminates is Fix A → rebuild → flash → cycle → fix B → rebuild → ... which serializes operator time onto rebuild latency.

Exceptions where a fix DOES require interim rebuild MUST be documented in the fix’s commit message as `REQUIRES_REBUILD: <reason>`. The default is the batch path; rebuild-now is the operator-authorized exception.

§11.4.10 — Credentials-handling mandate

Forensic anchor — direct user mandate:

“Credentials or any secret and sensitive data MUST NOT leak! This MUST BE added as the mandatory constraint into all Submodules and main project’s Constitution, CLAUDE.MD and AGENTS.MD.”

All credentials, secrets, API tokens, passwords, phone numbers, OAuth refresh tokens, signing keys, encryption keys, and any other authentication-or-authorization-bearing data used by test scripts, build automation, or any project tooling MUST be handled per the following rules without exception:

1. **No tracked-file storage.** Credentials NEVER live in any file that git tracks. Placeholder templates with `<replace ...>` or `EXAMPLE_VALUE_DO_NOT_USE` placeholders ARE allowed in `.example` files committed to show operators what keys to populate.
2. **Repository-wide ignore.** `.gitignore` MUST exclude all credential-bearing paths:
 - `.env`, `.env.*`, `*.env`, `.<service>.env`, `*.<service>.env`
 - `scripts/testing/secrets/*` with `.example + README.md` exception
3. **Runtime-load-only.** Tests load credentials AT EXECUTION TIME from operator-populated files under `scripts/testing/secrets/` (or per-project equivalent). If the file is missing, the test SKIPS with “credentials not configured for ” — it does NOT proceed without credentials and does NOT use defaults.
4. **No echo / no log.** Test scripts MUST NEVER print, log, or include credentials in error paths, stack traces, screen recordings, or output artifacts.
5. **Per-service file separation.** One file per service keeps blast radius small.
6. **Filesystem permissions.** `.env` files are `chmod 600`, parent directory is `chmod 700`.

7. **Rotation on suspected leak.** Rotate at provider, update local files, audit captured artifacts.

The consuming project's `.gitignore` MUST be verified before every commit: `git ls-files --cached | grep -E "\\..env$" MUST show no real .env files (only .env.example).`

§11.4.10.A — Pre-store credential leak audit (User mandate, 2026-05-17)

Forensic anchor — verbatim user mandate (2026-05-17):

“Us these for all future testing (full automation testing) and make sure they are not leaking anywhere or get git versioned!”

[Discovered during execution: the operator-provided test credentials ALREADY existed in 5 tracked files dating to prior project cycles. §11.4.10 + §11.4.30 catch new commits but did not detect the pre-existing leak when the operator re-provided the same values for a new use.]

When an operator provides credentials, API tokens, signing keys, or any other secret material to be stored in the project's gitignored configuration (`.env`, `scripts/testing/secrets/`, `secrets/`, etc.), the agent or human storing them MUST FIRST execute a repo-wide audit for prior leaks of THOSE specific values BEFORE storing.

The audit MUST:

1. **Grep every tracked file** for the literal credential value(s). `git ls-files | xargs grep -l <value>` is the canonical form. Search MUST cover all file types — `.md`, `.json`, `.yaml`, `.kt`, `.go`, `.py`, `.sh`, `.sql`, build configs, docs.
2. **Grep the entire git history** for the literal credential value(s). `git log -S<value> --all --source --remotes` reveals historical commits even if the current tree is clean. A history-only leak is a §11.4.10 violation of equal severity to a tree-leak — anyone who pulled an older commit still has the value in their working copy of any branch they cloned.
3. **Report findings to the operator BEFORE storing the new value in any operator-controlled location.** The operator may then decide whether to (a) rotate the credential at the upstream provider before reusing, (b) accept the historical leak as a known-compromise and proceed, or (c) abort the storage. Storing the value without surfacing the audit is itself a §11.4 PASS- bluff at the security layer.
4. **If a leak is found:** open a forensic incident record per the project's §6/§7 sixth-law-incidents discipline (or equivalent), redact the literal values from all tracked files in-place (replace with `<redacted-per-§11.4.10>` placeholders), and record an OPERATOR ACTION REQUIRED for rotation per §11.4.10 sub-clause 7. The redaction commit lives on master IMMEDIATELY; the history-purge follow-up (`git-filter-repo` / `BFG` / equivalent + force-push to every upstream) requires explicit per-operation operator authorization per §9.2.
5. **Strengthen pre-push hook detection** when a previously-undetected class of literal credential is discovered. The pre-push hook's credential-pattern grep MUST be extended to catch the specific pattern class that escaped (literal usernames matched with literal passwords in the same file, well-known credential SHA-256 hashes, org-specific naming conventions for service accounts, etc.). The extension MUST land in the same commit as the redaction.

Pre-build gate (recommended) CM-PRE-STORE-CREDENTIAL-AUDIT (when implemented in consuming project) checks the commit message body of any commit that touches `.env`, `.env.example`, `scripts/testing/` `secrets/`, or equivalent for an Audit-Pre-Stored: stamp citing the audit's outcome. Paired mutation (§1.1) strips the stamp and asserts gate FAILs.

Composes with §11.4.10 (the core credentials mandate this extends), §11.4.30 (`.gitignore` + `no-versioned-artifacts` — overlapping enforcement surface), §11.4.6 (`no-guessing` — audit findings stated as fact or as `PENDING_FORENSICS`), §11.4.7 (`demotion-evidence` — “I think it's not leaked” without `grep` evidence is a guess). Classification: universal (per §11.4.17). No escape hatch.

§11.4.11 — File-layout discipline

Project files **MUST** be organised by purpose, not by historical accident. Source code goes under canonical project roots. Tests go under canonical test directories. Logs and forensic artifacts go under operator-controlled directories (`~/Documents/`, `qa-results/`, `/tmp/`, OS-equivalent) — **NEVER** scattered at the repo root, **NEVER** inside working source trees, **NEVER** tracked unless they are reference assets explicitly required for evidence-replay.

Discovered drift is fixed by **moving** files to their canonical location and updating every caller — never by adding redirect shims or by leaving the misplaced copy “for backwards compatibility”.

Carve-out (User mandate 2026-05-20). The 5 canonical tracker documents — `docs/Issues.md`, `docs/Issues_Summary.md`, `docs/Fixed.md`, `docs/Fixed_Summary.md`, `docs/CONTINUATION.md` — sit at `docs/` root by design. They are architectural constants of the project layout, analogous to AOSP's `Makefile`, `Android.bp`, `OWNERS` files at repo root. Their location is encoded as literal path strings in §11.4.12 + §11.4.15 + §11.4.16 + §11.4.19 + §11.4.44 + §11.4.53 propagation gates plus the helper-script constellation that regenerates them (`generate_issues_summary.sh`, `generate_fixed_summary.sh`, `sync_issues_docs.sh`, `commit_docs.sh`, `colorize_progress_html.py`, `export_progress_docs.sh`). Moving them would require coordinated amendment of those 6 sister anchors plus 5 pre-build gates plus ~20 helper scripts plus 42 consumer files (parent + 10 owned submodules + nested + HelixQA dependencies) in a single PWU. Per §11.4.66, that scope is operator-blocked until explicitly authorised. Audit-snapshot files (`docs/audit/anti_bluff_audit.md`, `docs/audit/PRE_SONOS_TAG_READINESS.md`, `docs/audit/D1_WIFI_FAIL_CLASSIFICATION.md`, plus any future audit snapshots) **DO** move under `docs/audit/` per the §11.4.11 general principle — they are not part of the tracker constellation and carry no propagation-gate path-string burden.

§11.4.12 — Auto-generated docs sync mandate

Every auto-generated document (`Issues_Summary`, API reference, build manifest, etc.) **MUST** be regenerated in the **same commit** as any edit to its source. Stale auto-generated docs are §11.4 PASS-bluffs in document form: an operator reading them gets a divergent view of project state.

All three file types **MUST** stay in sync at all times — Markdown source + HTML export + PDF export (or equivalent multi-format output the project ships). Enforced by a pre-build gate that checks mtime ordering and content hash agreement.

§11.4.13 — Out-of-band sink-side captured-evidence

Whenever a downstream consumer (HDMI sink, cloud service, monitoring target, downstream system) provides a network-accessible introspection API that reports what was actually received, the test suite **MUST** consume that report as captured-evidence for every test asserting end-to-end delivery.

The on-source-side view **ALONE** is insufficient — that is the exact “tests pass but the feature doesn’t work” pattern §11.4 forbids. Sink-side reports **MUST** be: - Identity-verified (MAC, serial, UUID match) before consumed as evidence. - Topology-dispatched (§11.4.3) — sink-probe-unreachable → SKIP, never FAIL. - Cross-referenced with the on-source state at a matching wall-clock instant.

§11.4.14 — Test playback cleanup mandate

A test that completes (PASS, FAIL, or SKIP) **MUST** leave the target in a quiescent state — no orphan playback, no orphan recording, no orphan capture, no orphan background process continuing after the test’s last assertion. Tests are transient probes; their side- effects on shared target state **MUST NOT** outlive the test.

Mandatory protections:

1. Every test that issues a playback or capture command **MUST** issue the matching cleanup at test end (force-stop the app, stop the recorder, close the file descriptor, terminate the helper).
2. Cleanup is mandatory on **EVERY** exit path — PASS, FAIL, SKIP, error, trap, signal, parent-killed. Use shell trap '`<cleanup>`' EXIT or try/finally.
3. Tests **MUST** verify the cleanup succeeded via positive evidence (§11.4.5).
4. The orchestrator **MUST** run a post-test sanity check between tests and FAIL the just-completed test if it left orphan state. This blames the leaker, not the next-test victim.
5. No grace period for “the next test will clean it up” — that is precisely the §11.4 PASS-bluff pattern. Cleanup is the test’s responsibility, not the next test’s, not the orchestrator’s.

§11.4.15 — Item-status tracking mandate

Every active item tracked in the project’s Issues file **MUST** carry a ****Status:**** line within five lines of its heading. The status value is drawn from a closed set:

State	Meaning
Queued	In backlog, no work has started. Sub-state qualifier Queued – BLOCKED on <code><reason></code> permitted.
In progress	Active investigation or coding underway.
Ready for testing	

State	Meaning
	Source-side fix landed (pre-build / meta-test gates green), waiting for the next deploy cycle.
In testing	A deploy + cycle is in flight that exercises the fix; live captured-evidence is being collected per §11.4.2 + §11.4.5.
Reopened	Item failed its runtime test cycle and is back to active work. Carries a reference to the failure-cycle artefact.
Fixed (→ Fixed file)	Item closed: captured-evidence per §11.4.5 collected, runtime test PASSes, paired meta-test mutation FAILs, item migrated to the Fixed file.

Status MUST be updated as the item progresses. Every commit that changes an item's lifecycle state MUST update its Status line in the same commit.

All three Issues / Issues_Summary / Fixed file types MUST be in sync at all times (Markdown + HTML + PDF).

§11.4.16 — Item-type tracking mandate

Every active item tracked in the project's Issues file MUST carry a ****Type:**** line within eight non-blank lines of its heading (the same window the §11.4.15 Status audit uses). The type value is drawn from a CLOSED set of three values:

Type	Meaning
Bug	Product defect / regression / user-visible broken behaviour. The product worked before (or was expected to) and now does NOT for the end user.
Feature	New capability not previously offered to end users — a new integration, new output, new probe, new bank, new architectural surface.
Task	Internal workstream — not user-visible. Refactor, infrastructure, documentation, audit, covenant clause, gate, mutation, propagation enforcement, host-session-safety hardening. The largest class by count; the lowest-stakes default when the classification is ambiguous.

The vocabulary is CLOSED — only Bug, Feature, Task. Any other value is a violation. When ambiguous, fall back to Task.

Type tells the operator WHAT KIND of work an item is — distinct from severity (HOW URGENT) and status (WHERE IN LIFECYCLE). The three axes together drive ownership routing, testing strategy, release-note framing, and changelog classification.

The Issues_Summary file MUST carry the Type column for every active item. All three file types (Markdown + HTML + PDF) MUST stay in sync under the same CM-DOCS-EXPORT-SYNC discipline as §11.4.12 + §11.4.15.

Pre-build gates CM-ITEM-TYPE-TRACKING (scans Issues file for `**Type:**` presence within the audit window) and CM-COVENANT-114-16-PROPAGATION (anchor presence in every project CLAUDE.md / AGENTS.md) enforce the mandate. Paired mutations in the project’s meta-test prove neither gate is a bluff gate.

No escape hatch — items that genuinely don’t fit MUST be re-classified as INFORMATIONAL and excluded from the active-item set rather than carry a fabricated Type.

§11.4.17 — Universal-vs-project classification of new rules (User mandate, 2026-05-14)

Forensic anchor — direct user mandate (verbatim, 2026-05-14):

“Adding any new rules or mandatory constraints or anything relevant which should be added into Constitution, CLAUDE.MD, AGENTS.MD and relevant files MUST BE determined as reusable / universal VS project specific. If it is universal and reusable it MUST BE added into our root / main constitution Submodule, otherwise into our project level Constitution, AGENTS.MD and CLAUDE.MD (or Submodule if it is Submodule related).”

Before any new rule, mandatory constraint, covenant clause, gate declaration, or “MUST”-bearing statement is added to a project’s Constitution / CLAUDE.md / AGENTS.md or to a child submodule’s equivalents, the author MUST first classify it along this axis:

Classification	Definition	Destination
Universal (reusable)	The rule applies to ANY project — independent of language, framework, target platform, or domain. Examples: anti-bluff covenant, no-guessing mandate, credentials-handling, host-session safety, item-status tracking, multi-upstream push, file-layout discipline, deep-web-research-before-implementation.	This constitution submodule’s Constitution.md + CLAUDE.md + AGENTS.md. Propagates to every consuming project via inheritance.
Project-specific		

Classification	Definition	Destination
	The rule references a specific hardware target, vendor, model, SoC, vendor-fork, project-internal package name, deployment topology, or company-internal asset. Examples (illustrative — these are intentionally NOT in this universal file): the SoC’s specific power-management quirks; a particular Android/iOS/Linux SDK behaviour; a specific app or service bundled with the project.	The project’s own <code>Constitution.md</code> / <code>CLAUDE.md</code> / <code>AGENTS.md</code> (top-level), OR the affected submodule’s equivalents (when the rule scopes to that submodule).

A rule that mentions hardware part numbers, vendor names, specific package identifiers, geographic regions, or company-internal asset names is **project-specific by definition** — it cannot be lifted into the universal layer without genericising those references first. A universal rule MAY reference patterns (e.g., “USB-HID multitouch panels”) but MUST NOT name specific vendors.

Anti-bluff: the universal-vs-project classification is itself an auditable artefact. Every new rule’s commit message MUST include a one-line classification statement (e.g., `Classification: universal – applies to any project tracking items through an Issues/Fixed lifecycle`). A commit that adds a rule without classification justification is a §11.4.17 violation.

Authors who are uncertain SHOULD default to project-specific (narrower scope). Lifting a project-specific rule to universal later is cheap (one merge); generalising a universal rule that turned out to leak project-specific assumptions is expensive (every consumer must update).

Pre-build gate `CM-UNIVERSAL-VS-PROJECT-CLASSIFICATION` audits the last N commits for rule additions and asserts each one carries a classification statement. Paired mutation strips the classification literal and asserts the gate FAILs.

§11.4.20 — Subagent-driven-by-default mandate (User mandate, 2026-05-14)

Forensic anchor — direct user mandate (verbatim, 2026-05-14):

“Make sure we ALWAYS WORK EVERYTHING subagents-driven if it is possible or applicable!”

When operating in a multi-agent runtime (Claude Code with subagents, Cursor with task-runners, Aider with sub-sessions, or any CLI tool that supports delegated execution), the primary agent **MUST** default to subagent delegation for any task that satisfies **AT LEAST ONE** of:

1. **Multi-step scope** — three or more discrete editing / file-creation / verification phases. Hand-off boundaries are where stalls happen; pre-planned subagent decomposition avoids them.
2. **Parallelisable work** — two or more independent tasks with no shared mutable state. Parallel subagents finish wall-clock-faster than sequential foreground work **AND** insulate the primary agent's context window from the volume of file reads.
3. **Long-running diagnostic loops** — repeated probe / re-test cycles, build-flash-cycle loops, soak tests. Subagents survive context compression boundaries that would otherwise truncate progress.
4. **Domain-specific or specialised workflows** — code review, security audit, infrastructure change review, performance triage, documentation propagation across N files. Specialised subagents (code-reviewer, iac-reviewer, general-purpose) bring pre-curated tool sets and disciplines the primary agent would re-derive from scratch.

The primary agent **SHOULD** only do foreground work when **AT LEAST ONE** of:

- The task is a single edit or single file read with no follow-up.
- The task requires conversational clarification with the operator mid-execution (subagents cannot ask the operator questions).
- The task is gating critical state (a commit, a push, a tag cascade) that must be sequenced before the next operator action.
- The task is so quick (under ~30 seconds of execution) that subagent spin-up overhead exceeds the work itself.

Anti-stall discipline. Subagents have watchdog timeouts in most runtimes (Claude Code: ~600 s of no-progress). When delegating, the parent agent **MUST**: (a) scope tightly (no “complete everything” prompts — break into 4-6 tightly-scoped tasks); (b) require **checkpoint commits** after each major task so partial progress is preserved even on stall; (c) explicitly prohibit destructive operations the subagent might attempt to “be helpful” (`git reset --hard`, `git push --force, --no-verify`); (d) provide explicit discipline pointers (§11.4.6 no-guessing, §11.4.10 credentials, §11.4.17 classification).

Anti-bluff applied. A subagent that returns “completed” without captured evidence in commits / outputs / artifacts is a §11.4 PASS-bluff at the multi-agent layer. The parent agent **MUST** verify subagent claims against repo state (`git log`, `git status`, post-completion gate runs) before treating the work as landed.

Coordination. When parallel subagents touch overlapping files, the parent agent **MUST** partition the work so subagents work on **non-overlapping files**. The parent's `commit_all.sh --auto-cascade` naturally bundles both subagents' dirty files into atomic commits via `git add -A` — exploit this for “forward-reference + provider” patterns where one subagent creates files another subagent references.

No escape hatch. Operating exclusively in the foreground when subagent delegation is feasible burns operator wall-clock time, risks context-window overflow on large tasks, and forfeits the parallelism that multi-agent runtimes were designed to provide. Pre-build gate `CM-SUBAGENT-DELEGATION-AUDIT` (when

implemented per consuming project) scans recent session transcripts for multi-step foreground work that should have been delegated; paired mutation enforces the gate is not a bluff.

§11.4.18 — Script documentation mandate (User mandate, 2026-05-14)

Forensic anchor — direct user mandate (verbatim, 2026-05-14):

“Make sure that every single script inside the scripts dir (in all subdirs in depth) is fully documented and covered with full user manuals and user guides! ... Whenever is some bash script modified we MUST document it fully and create for it complete user guide(s) and manual(s) - if already exists make sure all of it is in sync and fully updated! No documentation ever can be out of sync with its codebase!”

Every Bash / shell / POSIX-sh script ANYWHERE in a project (anywhere under scripts/, bin/, tests/, library directories, Upstreams/, deployment hooks, CI helpers, etc. — depth-N recursive) MUST carry:

1. **In-source documentation block** at the top of the file:
 - Purpose: one-sentence description of what the script does.
 - Usage: complete invocation syntax with all flags, env vars, positional arguments, examples.
 - Inputs: env vars, files read, command-line args, stdin.
 - Outputs: files written, stdout/stderr behaviour, exit codes.
 - Side-effects: anything that changes system state outside the script’s own scratch directory.
 - Dependencies: required commands (and how to install them) + required system state (e.g., “must be run as root”, “must run from project root”).
 - Cross-references: companion guides under docs/.
2. **External user guide** under docs/scripts/<script-name>.md (Markdown).
Required sections:
 - **Overview** — what the script does in plain English.
 - **Prerequisites** — environment + dependencies.
 - **Usage examples** — copy-paste recipes for the common cases.
 - **Edge cases** — unusual invocations, error conditions, how to diagnose them.
 - **Internal behaviour** — for advanced users / future maintainers: what the script does step-by-step.
 - **Related scripts** — companion scripts and how they compose.
 - **Last verified version / date** — when the doc was last reconciled against the script source.

Whenever the script source is modified, the in-source documentation block AND the external user guide MUST be updated in the SAME commit. A commit that touches a script but leaves its documentation unchanged (or worse, lets it drift) is a §11.4.18 violation.

Pre-build gate CM-SCRIPT-DOCS-SYNC walks every *.sh and *.bash under scripts/ (or equivalent script directories), verifies a companion docs/scripts/<name>.md exists, AND verifies the doc was modified in the same commit as the script (by SHA cross-reference, or by mtime ≥ script mtime as a softer floor). Paired mutation strips the doc-sync invariant and asserts the gate FAILs.

No documentation ever can be out of sync with its codebase.

This mandate composes with §11.4.11 (file-layout discipline — docs live under docs/, scripts live under scripts/) and §11.4.12 (auto-generated docs sync — the on-disk Markdown + its rendered HTML/PDF must all stay synchronised).

§11.4.19 — Fixed-document column-alignment mandate (User mandate, 2026-05-14)

Forensic anchor — direct user mandate (verbatim, 2026-05-14):

“Make sure that Fixed document (all 3 file formats) is consistent and has the same structure and columns as Issues and Issues_Summary documents! This has to be added as detail to Constitution, CLAUDE.MD and AGENTS.MD (constitution Submodule where these information shall exist).”

Every project that maintains an Issues / open-work tracker AND a Fixed / closed-archive tracker MUST keep the two structurally aligned along the same lifecycle axes — at minimum **Status** and **Type** — so an operator reading either gets a consistent view across the entire lifecycle (open → in progress → ready for testing → in testing → reopened → fixed).

Specifically:

1. **Per-entry annotation.** Every heading in the Fixed archive document MUST carry, within 8 non-blank lines of the heading, a ****Status:**** line and a ****Type:**** line — exactly as §11.4.15 + §11.4.16 require for the Issues tracker. For Fixed entries the Status value is drawn from the **closure-set**:

- Fixed (→ Fixed.md) — fully verified on device with captured evidence per §11.4.5.
- Fixed – pending device verification — source-side fix landed, regression-protection gate wired, awaiting the next firmware build + flash + cycle for captured evidence.
- Fixed – RECLASSIFIED — item was reclassified during investigation (test-side defect, environment-only, etc.).

Type values follow §11.4.16: Bug | Feature | Task. Default on missing-annotation: Task.

2. **Auto-generated companion summary.** A Fixed_Summary.md MUST exist alongside Fixed.md, regenerated by an automation script (e.g. generate_fixed_summary.sh), and MUST mirror the column structure of Issues_Summary.md exactly:

```
| # | Level | Status | Type | One-line description |
```

The Level (severity) column uses the **original severity** the item had when it was open — so an operator can correlate a closed item with its original priority bucket. The same deterministic classification rules apply (REAL PRODUCT DEFECT / PARTIAL-or-WORKSTREAM-or-DEEP-DIVE / everything-else mapping to C / M / L).

3. **Three-format sync.** All three file types (.md, .html, .pdf) for BOTH the Fixed archive AND the Fixed_Summary MUST stay in sync at all times — same .html + .pdf export pipeline that §11.4.12 requires for Issues + Issues_Summary. The single-shot sync wrapper (e.g. sync_issues_docs.sh) MUST invoke the Fixed_Summary generator AND export all five doc variants (Issues, Issues_Summary, Fixed, Fixed_Summary, CONTINUATION) in one operation so they all carry the same mtime.
4. **Cross-lifecycle consistency.** A status line in Issues.md that reads Fixed (→ Fixed.md) is a **migration marker**, not a self-describing closure: when an Issues.md entry resolves, it MUST be moved to Fixed.md in the same commit, then disappear from Issues_Summary (because Issues_Summary only enumerates OPEN items) and appear in Fixed_Summary (because Fixed_Summary enumerates CLOSED items). Per §11.4.4 fix-closure protocol + §11.4.5 captured-evidence requirement, the migrated entry in Fixed.md MUST carry: closure cycle, closure commit SHA, captured evidence pointer, regression-protection gate name + paired mutation reference.

Pre-build gate CM-FIXED-COLUMN-ALIGNMENT (5+ invariants): - docs/Fixed_Summary.md exists. - docs/Fixed_Summary.md table header line contains both Status and Type columns (column-alignment with Issues_Summary). - mtime(Fixed_Summary.md) >= mtime(Fixed.md) (regenerated after Fixed.md edits). - Generator script exists (generate_fixed_summary.sh or equivalent integrated branch in generate_issues_summary.sh). - Sync wrapper invokes the Fixed_Summary generator (grep for the script name in sync_issues_docs.sh or equivalent). - HTML + PDF exports for both Fixed and Fixed_Summary exist (docs/Fixed.html, docs/Fixed.pdf, docs/Fixed_Summary.html, docs/Fixed_Summary.pdf).

Paired mutation: strip the Status column from Fixed_Summary.md table header (or rename the generator) → gate FAILs. Enforced by meta_test_false_positive_proof.sh.

Propagation. This anchor is a §11.4.17-classified **universal** rule — it composes naturally with §11.4.12 (auto-generated docs sync), §11.4.15 (Status tracking), §11.4.16 (Type tracking) and is mandatory across every consuming project's CLAUDE.md / AGENTS.md. The propagation gate CM-COVENANT-114-19-PROPAGATION (when implemented per consuming project) enforces the anchor's presence in every covenant file.

No escape hatch. A Fixed archive that lacks Status/Type columns or whose Summary companion drifts out of sync is the exact “different lifecycles report different facts” hazard §11.4 forbids — an operator reading Issues_Summary can no longer answer “where is item X now?” by cross-referencing Fixed_Summary because the schemas diverge. Non-compliance is a release blocker regardless of context.

§11.4.21 — Operator-blocked status + self-resolution exhaustion mandate (User mandate, 2026-05-14)

Forensic anchor — direct user mandate (verbatim, 2026-05-14):

“Add into Issues and Issues_Summary new status: Operator blocked! We need another column with details in Issues document with exact details on how exactly is operator blocked issue! Before issue becomes operator blocked we MUST investigate in-depth all ways on how it can be resolved without operator’s involvement - by you, or by activating any relevant mechanism for the resolution!”

Every project that maintains an Issues / open-work tracker MUST treat operator dependency as a **last-resort classification**, earned only after documented exhaustion of self-resolution paths. Routing an item to the operator without proving the agent + tooling cannot unblock it themselves is the §11.4 PASS-bluff pattern transplanted into the planning layer — the work stalls indefinitely while the agent claims “nothing more I can do.”

1. Status vocabulary extension. §11.4.15’s closed-set of Status values is extended with a seventh value, Operator-blocked. The full closed-set becomes:

#	Status value	Meaning
1	Queued	Awaiting work to start.
2	In progress	Active agent / developer work.
3	Ready for testing	Source-side fix landed, awaiting verification.
4	In testing	Captured-evidence cycle running.
5	Reopened	Previously closed item resurfaced (regression).
6	Operator-blocked	NEW. Cannot proceed without an operator action that the agent + available tooling cannot perform.
7	Fixed (→ Fixed.md)	Closure migration marker per §11.4.15 + §11.4.19.

2. Self-resolution exhaustion mandate. BEFORE classifying any item as Operator-blocked, the agent MUST verifiably exhaust every self-resolution path applicable to the project:

- **(a) CLI / ADB / SSH / API access** — can the work be done via any access the agent already has? (e.g. `adb shell pm clear, gh api, aws ssm, kubectl exec, psql`).
- **(b) Subagent delegation** — can a specialised subagent (per §11.4.20) reach further than the primary agent? (e.g. general-purpose for multi-step probing, code-reviewer for architectural blockers, iac-reviewer for infra access).
- **(c) Existing tooling** — is there a script / helper / library already in the repo that handles this case? (e.g. `scripts/testing/<helper>.sh`, project-specific automation).
- **(d) Captured fallback** — can the blocker be sidestepped with a synthetic event, test asset substitution, mock, or topology-aware SKIP per §11.4.3?

(e.g. simulate the input, substitute a stand-in fixture, downgrade to a pre-build gate).

- **(e) Documentation + research** — per §11.4.8, has the agent consulted external sources (official docs, vendor guides, open-source codebases, issue trackers) for a self-resolution pattern the community has already solved?

Only AFTER each applicable path is **verifiably attempted and documented as exhausted** is Operator-blocked the correct classification. “I didn’t try X because I assumed it wouldn’t work” is a §11.4.6 no-guessing violation AND a §11.4.21 self-resolution violation.

3. Operator-Block-Details mandate. Every Status: Operator-blocked item MUST carry an ****Operator-Block-Details:**** line within 8 non-blank lines of its heading (mirroring the §11.4.15 Status placement pattern). The content MUST state ALL of:

- **WHAT** — the specific concrete action the operator must perform (verb + object + target system). Generic “the operator must investigate” is forbidden — name the action.
- **WHY** — the alternatives that were exhausted, enumerated. Each exhausted path from §11.4.21.2 above gets a one-line statement of what was attempted and why it could not unblock the item. “Subagent delegation: tried general-purpose with X scope — blocked because Y” is the bar.
- **UNBLOCK CONDITION** — the observable signal that the operator has completed the action (e.g. “operator confirms hardware rev-B installed”, “operator pastes new API key into scripts/testing/secrets/.foo.env”, “operator force-pushes PR #123 merge to upstream”). Without this, the agent cannot automatically detect when to re-evaluate.
- **WHO** — the handle / contact / pointer to the document where operator-side details live if questions arise (e.g. “operator: @username”, “see docs/operator/<topic>.md”, “vendor: support-id 12345”). Operator-blocked without WHO is operationally untrackable.

4. Issues_Summary inclusion as sortable axis. The auto-generated `Issues_Summary.md` (per §11.4.12 + §11.4.15) MUST include Operator-blocked as a first-class Status value in the table — the column header is unchanged but operators can sort / filter on it deterministically. The summary generator MUST also emit a count of Operator-blocked items in any rollup line / header.

5. Periodic re-evaluation requirement. Items in Operator-blocked status MUST be re-evaluated every Nth tag-cycle (project-defined, recommended every 3rd tag cycle). Operator dependencies change over time: CI may have been added, hardware may have been swapped, automation may have matured, vendor APIs may have unlocked. An item that was correctly Operator-blocked at tag T may be self-resolvable at tag T+3. Re-evaluation MUST follow the same §11.4.21.2 exhaustion checklist and document each re-attempt.

6. Anti-bluff layer. A fake Operator-blocked (classification applied without exhausting §11.4.21.2 alternatives) is a §11.4 covenant violation at the planning layer — equivalent severity to a PASS-bluff at the testing layer. The agent’s commit message introducing or maintaining an Operator-blocked classification MUST contain evidence of self-resolution attempts (the explicit “Attempted: a — exhausted because X; b — exhausted because Y; c — exhausted because Z” form).

7. Pre-build gates (recommended, per consuming project):

- **CM-ITEM-OPERATOR-BLOCKED-DETAILS** — every heading whose Status line equals Operator-blocked has an ****Operator-Block-Details:**** line within 8 non-blank lines. Paired mutation strips the details line → gate FAILs.
- **CM-OPERATOR-BLOCKED-SELF-RESOLUTION-AUDIT** — every NEW Operator-blocked item introduced in a commit carries an “Attempted: ...” audit trail (either in the Issues.md entry body OR in the commit message). Paired mutation introduces an Operator-blocked item without the audit trail → gate FAILs.

Propagation. This anchor is a §11.4.17-classified **universal** rule — it composes with §11.4.15 (Status tracking), §11.4.16 (Type tracking), §11.4.19 (Fixed-document column alignment), §11.4.12 (auto-generated docs sync), §11.4.20 (subagent delegation as a self-resolution path), §11.4.8 (research-before-implementation as a self-resolution path), §11.4.6 (no guessing about what the operator “would have to do”). Propagation gate CM-COVENANT-114-21-PROPAGATION (when implemented per consuming project) enforces the anchor’s presence in every CLAUDE.md / AGENTS.md across the parent + every owned submodule + every dependency.

No escape hatch. Items that legitimately need operator intervention are real and unavoidable — hardware swaps, vendor contracts, physical-world inputs, account credentials, irreversible business decisions. But the classification must be **earned** via documented exhaustion. An Operator-blocked heading without an audit trail of self-resolution attempts is a §11.4 release blocker regardless of context.

§11.4.22 — Document-sync commit discipline (User mandate, 2026-05-14)

Forensic anchor — direct user mandate (verbatim, 2026-05-14):

“For Issues, Issues_Summary and Fixed docs (and all its exports) we MUST commit and push (only them) as soon as they are updated (synced). Continuation doc and other similar / relevant documentation MUST be committed and pushed too! Like this we will always have up to date status of working items without need to do full commit and push of everything if that is not possible to do!”

Classification: §11.4.17-classified **universal** — every project that tracks work items through an Issues / Fixed lifecycle has the same problem, so the discipline lives in the Constitution and every consuming project implements its own wrapper.

The defect this anchor closes. Until this anchor existed, the only way to push a status update was the full-tree commit wrapper. When the working tree had unrelated heavy churn (large submodule diffs, in-flight rebase, partial-network conditions, ongoing build), operators had two bad options: (a) wait until the heavy work finishes — which can be hours — leaving the doc-status stale to every other agent; (b) skip the doc-status update entirely — which is a §11.4.4 documentation-drift violation. Neither option is acceptable.

The mandate. Every project under this Constitution that ships a status-tracking doc set MUST provide a **lightweight commit path** distinct from the full-repo commit wrapper. The lightweight path stages, commits, and pushes **only** the status-tracking doc set:

1. The active issue tracker (docs/Issues.md in ATMOSphere's layout; the project's equivalent file elsewhere) and its HTML + PDF exports.
2. The auto-generated issues summary (docs/Issues_Summary.md) and its HTML + PDF exports.
3. The fixed-bug archive (docs/Fixed.md) and its HTML + PDF exports.
4. The auto-generated fixed summary (docs/Fixed_Summary.md) and its HTML + PDF exports.
5. The cross-session continuation document (docs/CONTINUATION.md, per §12.10) and its HTML + PDF exports.
6. Any auto-generated audit artifact (docs/audit/anti_bluff_audit.md or equivalent) that pairs with the above.

The lightweight wrapper MUST:

1. **Auto-invoke the project's export-regeneration pipeline first** — so Markdown + HTML + PDF are all synchronised before push. (Skipped only when an explicit `--no-sync` flag is passed AND the operator just ran the pipeline manually.)
2. **Stage ONLY the doc-set files** — `git add` of an explicit file list, NEVER `git add -A` (which would catch unrelated WIP churn).
3. **Use a separate flock** disjoint from the full-tree commit wrapper's lock — so the two can run in parallel without contention on disjoint file sets.
4. **Push to every parent-repo remote** (reusing the project's push-all driver where available).
5. **Exit code semantics** — 0 success, 1 validation failure, 2 lock contention, 3 nothing-to-commit (informational, not an error — so the wrapper can be wired into automation that tolerates no-op outcomes).

The lightweight wrapper MUST be invocable either standalone OR via a flag on the full-tree wrapper (e.g. `commit_all.sh --docs-only`) so operators have a single mental model.

§9 preflight inheritance. The lightweight wrapper inherits the parent project's §9 preflight discipline — if `meta_test_*` is mid- mutation or `*.mut.bak` files exist, the wrapper refuses to run, same as the full-tree wrapper.

Pre-build gates (recommended, per consuming project):

- **CM-COMMIT-DOCS-EXISTS** — verifies the lightweight wrapper exists + is executable, its external user guide (per §11.4.18) exists, the full-tree wrapper advertises the delegation flag, and the wrapper's doc-set enumeration lists $\geq N$ entries (N depends on project — ATMOSphere's set is 15). Paired mutation strips the doc-set array → gate FAILs.

Propagation. Composes with §11.4.12 (export-sync invariant — HTML + PDF in sync with Markdown after every edit), §11.4.15 (item status tracking — Status lines must be visible quickly), §11.4.18 (every script ships with an in-source doc block + an external user guide), §12.10 (CONTINUATION document maintenance — the lightweight wrapper is one of the mechanisms that keeps CONTINUATION current).

No escape hatch. The discipline exists because doc-status drift is itself a §11.4 PASS-bluff pattern at the documentation layer: a green-looking `Issues_Summary.html` that doesn't reflect Markdown reality is the same class of defect as a passing test that doesn't reflect runtime reality. Operators who can't run the lightweight wrapper for some reason **MUST** run the full-tree wrapper and accept the heavier cost — the doc-status drift is non-negotiable.

§11.4.23 — Visual-cue & grouping mandate for Issues docs (User mandate, 2026-05-14)

Forensic anchor — direct user mandate (verbatim, 2026-05-14):

“We **MUST** make sure that Issues, Issues_Summary and Fixed docs and its exported files (PDFs and HTMLs) have one small improvement: we **MUST** introduce some background coloring for cells of item type and status! Also, we **MUST** group items by the status! ... Base colors for types would be: bug - background pale red, task - background pale blue, feature - background pale yellow. However, changing the status affects the color of the cell of the type and the status cell background color: queued - no color, no effect, fixed - both cells pale green, in progress - only status cell is pale green, reopened - status cell is pale red, blocker - status cell is red (not pale, however **MUST BE** still readable), anything else please follow the same logic!”

Classification: §11.4.17-classified **mixed** — the discipline (visual cues + status grouping for tracked-item docs) is universal across every project that maintains an Issues / Fixed lifecycle. The implementation (CSS classes, HTML post-processor, weasyprint integration) is project-specific because each project's export toolchain differs (pandoc, asciidoctor, sphinx, etc.).

The defect this anchor closes. A long, multi-column tracked-item table where every row looks identical at first glance is a §11.4 operational-readability defect — operators have to *read* every Status column to identify what is queued vs in-progress vs fixed. With dozens to hundreds of items, this is slow enough that operators inevitably skim, and the doc effectively becomes write-only. Color cues + status grouping make at-a-glance triage possible.

The mandate. Every project under this Constitution that ships tracked-item docs (per §11.4.15) **MUST** apply visual-cue coloring AND status grouping to those docs' HTML + PDF exports. The Markdown source stays free of color noise (Markdown is the canonical text), but the HTML + PDF exports **MUST** be enriched in a post-processing stage:

Type cell base colors:

Type	Background
Bug	pale red (#FFE5E5 or equivalent — ~5% red saturation)
Task	pale blue (#E5F0FF)
Feature	pale yellow (#FFF8DC)

Status cell colors (override the Status column; Fixed also overrides Type cell):

Status	Status cell	Effect on Type cell
Queued	no color (transparent)	keeps base type color
In progress	pale green (#D4F1D4)	keeps base type color
Ready for testing	pale green-blue (#C6EAEF)	keeps base type color
In testing	pale green-yellow (#E7F4D6)	keeps base type color
Reopened	pale red (#FFCCCC)	keeps base type color
Fixed	pale green (#A8E6A8)	also pale green
Operator-blocked / Blocker	vibrant red (#FF4444) with white text for readability	keeps base type color

Status grouping: items in the rendered table **MUST** be grouped by Status, emitting an H3-class heading for each Status group with the item count. Group order (most-actionable first; closed last):

1. Operator-blocked / Blocker
2. In testing
3. Ready for testing
4. In progress
5. Reopened
6. Queued
7. Fixed

This ordering surfaces the items requiring operator attention at the top of the doc; closed/archived items sink to the bottom.

Print-fidelity requirement. The CSS **MUST** declare `print-color-adjust: exact` (or the project's renderer's equivalent) so PDF exports preserve the cell backgrounds. A weasyprint default-rendered PDF that strips colors is a §11.4 export-drift defect — HTML and PDF **MUST** be visually equivalent.

Pre-build gates (recommended, per consuming project):

- **CM-DOC-COLOR-GROUPING-DISCIPLINE** — verifies the project's CSS carries the `type+status` class set, the post-processor exists and is executable, the sync wrapper invokes it, AND after sync the rendered HTML carries the grouping-heading + per-cell color classes (the captured-evidence proof — without this last invariant a stub post-processor could silently pass the static gate). Paired mutation strips a Status class literal from CSS → gate FAILs.

Propagation. Composes with §11.4.12 (export-sync invariant — HTML + PDF in sync with Markdown after every edit), §11.4.15 (item status tracking — every active item carries a Status line from the closed-set vocabulary), §11.4.19 (Fixed_Summary column-aligned companion), §11.4.22 (lightweight commit path for doc-set updates).

No escape hatch. Color cues are not optional polish — they are the operational-readability seam that converts a long tracked-item table from a write-only Markdown blob into a self-triaging document.

§11.4.24 — Build-resource stats tracking mandate (User mandate, 2026-05-14)

Forensic anchor — direct user mandate (verbatim, 2026-05-14):

“We need to incorporate through the Containers Submodule we use to run our System build Container the following mechanism safely: during its working lifetime we **MUST** track in proper Markdown document exported into PDF and HTML stats on System resources use! What was the peak and the minimum **EVER** the Container / System building process used! ... In top of the document we **MUST** have this ever values. ... After this we **MUST** present properly sorted divided by version tags names all build processes we were running with notes if the process completed with success or not, and resources usage during the process! Besides min and max values ever and per build iteration we **MUST** have for each resource we have average values — average memory usage, average CPU usage, and all other parameters IO and other resources!”

Classification: §11.4.17-classified **mixed** — the discipline of tracking host-side build-resource usage with min / max / mean / p95 per-build PLUS ever-values across all builds in a versioned Markdown + HTML + PDF triple is universal across every long-build project (AOSP, kernel, large monorepo, ML model training, multi-hour data pipelines). The implementation (registry path, monitor script name, exporter wiring) is project-specific.

The defect this anchor closes. Long builds (multi-hour AOSP builds, ML training runs, large monorepo CI) consume highly variable resource profiles over their lifetime — soong_build’s 33 GB peak coexists with quiescent linker-only stretches at <2 GB. Without per-build sampling + ever-values, operators have NO empirical basis to (a) size containers correctly, (b) detect resource regressions across versions, (c) prove which build iteration was the OOM-killer’s trigger, (d) reason about parallelism caps (cf. §12.7 -j2 hard cap whose forensic basis would have been one full Stats.md generation earlier had this discipline been in place). Memory pressure debugging without time-series data is the bluff this anchor forbids.

The mandate. Every project under this Constitution with a build exceeding **1 minute wall-clock** **MUST**:

1. Run a host-side resource sampler for every build. The sampler runs as a sibling background process (not inside the build’s own process tree), reads /proc/meminfo + /proc/loadavg + /proc/stat + /proc/diskstats (or platform equivalents on non-Linux) at a fixed interval (recommended 5 s default — fine-grained enough to catch seconds-long peaks; coarse enough to keep itself <5% CPU and <50 MB RSS). Samples written as TSV.

2. Compute per-build summaries on stop. Per metric (memory used, CPU%, load average, disk read/write): **min, max, mean, p95**. p95 is the 95th-percentile sample value — required because mean alone hides extended-peak behaviour (a 2 h build with one 33 GB soong_build peak has mean ≈ 18 GB but p95 ≈ 32 GB; the latter is what operators must size for).

3. Append per-build summaries to a registry. One TSV row per build (build_id, status SUCCESS/FAIL/UNKNOWN, start/end timestamps, per-metric min/max/mean/p95, total disk read/write). The registry is the single source of truth — the Markdown report is derived.

4. Maintain ever-values across ALL builds. Min / max / mean across every tracked build, computed on every Markdown regeneration from the registry. Surface at the **top** of the report per User mandate verbatim.

5. Markdown + HTML + PDF triple. Stats.md (auto-generated), exported to Stats.html + Stats.pdf through the project's normal export pipeline. Triple stays in sync per §11.4.12. Committed via the project's lightweight doc-sync wrapper per §11.4.22.

6. Sort per-build entries by recency (most recent first) AND, where applicable, group by version tag. Operators reading the report want the latest build's profile at the top and the historical sweep underneath.

7. Track status per build. SUCCESS / FAIL / UNKNOWN — and if FAIL, capture the reason (exit code, OOM, ANR, etc.). A FAIL build's resource profile is the most operationally interesting one in the report — never silently drop FAIL entries.

8. Performance budget for the sampler itself. <50 MB RSS, <5% CPU. The monitor **MUST NOT** itself contribute to the resource pressure it is measuring (Heisenberg-class observer effect at scale). Pure /proc reads + awk arithmetic is the reference implementation; vmstat / pidstat acceptable; psrecord / heavy-Python-import samplers are forbidden.

9. Survive build failures. The stop hook **MUST** be called from both the success AND failure paths of the build wrapper so FAIL builds still produce a row. Hooking only the success path is a PASS-bluff at the telemetry layer.

Pre-build gate (recommended, per consuming project):

- **CM-BUILD-RESOURCE-STATS-TRACKER** — verifies (a) the monitor script exists + executable, (b) the Stats.md target file exists,
 1. the build wrapper invokes start AND stop, (d) the doc-sync wrapper enumerates the Stats.{md,html,pdf} triple, (e) the exporter advertises the build-stats slug. Paired mutation hides the monitor script aside → gate FAILs.

Propagation. Composes with §11.4.12 (export-sync invariant — HTML + PDF in sync with Markdown), §11.4.18 (script-documentation discipline — the monitor ships with an external user guide), §11.4.22 (lightweight doc-sync commit wrapper enumerates the Stats.{md,html,pdf} triple), §12.6 / §12.7 / §12.9 (host-safety + containerized-build envelope whose forensic anchors are the empirical motivation for this telemetry).

No escape hatch. Build-resource tracking is not optional polish — it is the operational-evidence seam that converts post-mortem “the build OOM'd, no idea why” into “build X at git-SHA Y peaked at Z GB at minute N, which is W% above the rolling p95”. The discipline pays for itself the first time it diagnoses a build-resource regression.

§11.4.25 — Full-Automation-Coverage Mandate (User mandate, 2026-05-15)

Forensic anchor — verbatim user mandate (2026-05-15):

“Make sure that every feature, every functionality, every flow, every use case, every edge case, every service or application, on every platform we support is covered with full automation tests which will confirm anti-bluff policy and provide the proof of fully working capabilities, working implementation as expected, no issues, no bugs, fully documented, tests covered! Nothing less than this does not give us a chance to deliver stable product! This is mandatory constraint which MUST BE respected without ignoring, skipping, slacking or forgetting it!”

Operative rule. For every consuming project under this Constitution, no feature, functionality, flow, use case, edge case, service, or application on any supported platform may be considered **deliverable** until it is covered by automation tests that collectively prove six invariants:

1. **Anti-bluff posture (per §7.1 + §11.4):** every assertion carries captured runtime evidence; metadata-only / configuration-only / grep-based / absence-of-error PASSes are forbidden.
2. **Proof of working capability:** the user-visible behaviour is exercised end-to-end on the target platform topology (per §11.4.3), not in a mock or in-memory facsimile.
3. **Working implementation as expected:** assertions match the product’s documented promise (in user manual, README, specs), not an implementation-detail back-door.
4. **No issues, no bugs:** the test suite is the canonical seam for surfacing defects — passing without producing defect signals means defects do not exist or have been previously surfaced, tracked (per §11.4.15 / §11.4.16), and closed.
5. **Fully documented:** the feature has a user-facing doc entry (per §11.4.18 for scripts, and project-level user-manual coverage for everything else); the doc is kept in sync with the tests (per §11.4.12).
6. **Tests-covered (the four-layer floor per §1):** pre-build presence-of-change gate, post-build artifact-shipped gate, runtime / integration / on-device gate, AND meta-test paired mutation that proves the runtime gate catches the break.

Cross-cutting reach. This mandate is **universal** — it applies to every project consuming this Constitution, every feature surface they expose (HTTP endpoints, CLI commands, slash commands, IPC channels, plugins, hooks, MCP servers, agents, providers, integrations, GUI flows, mobile flows, scheduled jobs), and every supported platform (Linux, macOS, Windows, iOS, Android, AuroraOS, HarmonyOS, embedded, containers, headless servers, kiosk, etc.). A project that ships a feature without satisfying all six invariants is **not delivering a stable product**, irrespective of how green its summary line looks.

Coverage audit. Consuming projects MUST publish a coverage ledger (matrix of: feature × platform × invariant-1..6 × status) that is regenerated as part of the release-gate sweep. The ledger itself is documented (per §11.4.18 / §11.4.12) and committed via the §11.4.22 lightweight doc-sync wrapper. Gaps in the ledger (**UNCONFIRMED:** / **PENDING:** / **BLOCKED:** cells) MUST cite a tracked work item per §11.4.15 + §11.4.16; rows that quietly omit a platform are §11.4.25 violations.

Composition. §11.4.25 explicitly stacks on top of §1 (four-layer test-coverage floor), §1.1 (false-positive immunity), §7.1 (positive-evidence-only validation), §11.4.1 (FAIL-bluffs forbidden), §11.4.2 (recorded-evidence requirement), §11.4.3 (per-environment-topology dispatch), §11.4.6 (no-guessing), §11.4.15 / §11.4.16 (status + type tracking), §11.4.17 (universal-vs-project classification — this rule is

universal), §11.4.18 (script documentation), §11.4.20 (subagent delegation when the work is multi-step), §11.4.22 (lightweight doc-sync). It does NOT supersede them; it forecloses the loophole “the feature works locally for me, ship it”.

Classification: universal (per §11.4.17). No escape hatch. Severity-equivalent to a §11.4 PASS-bluff at the release-gate layer — a project claiming “Done” without honoring §11.4.25 is making a false claim regardless of how the work-item tracker labels the row.

§11.4.26 — Constitution-Submodule Update Workflow Mandate (User mandate, 2026-05-15)

Forensic anchor — verbatim user mandate (2026-05-15):

“Every time we add something into our root (constitution Submodule) Constitution, CLAUDE.MD and AGENTS.MD we MUST FIRST fetch and pull all new changes / work from constitution Submodule first! All changes we apply MUST BE committed and pushed to all constitution Submodule upstreams! In case of conflict, IT MUST BE carefully resolved! Nothing can be broken, made faulty, corrupted or unusable! After merging full validation and verification MUST BE done!”

Operative rule. Before ANY agent or operator modifies the constitution/Constitution.md, constitution/CLAUDE.md, or constitution/AGENTS.md files of a project that consumes this Constitution as a submodule, the agent or operator MUST execute the following pipeline in order, with NO step skipped:

1. **Fetch + pull first.** From inside the constitution/ submodule worktree, run `git fetch <every-configured-remote>` followed by `git pull --ff-only origin <branch>` (or `--rebase` if non-FF-mergeable; **never** `--strategy=ours` / `--allow-unrelated-histories` without explicit operator authorization). The submodule MUST be at upstream tip BEFORE any local edit is applied.
2. **Apply the change.** Edit the relevant file(s). The edit MUST classify itself per §11.4.17 (universal vs project-specific) — only universal additions belong in the constitution submodule; project-specific clauses belong in the consuming project’s own governance files. The edit MUST include the verbatim user mandate (if it originated from one) as a forensic anchor.
3. **Validate before commit.** Run the constitution submodule’s `meta_test_inheritance.sh` (or equivalent) to confirm the inheritance chain still resolves. Verify no governance file was left with merge-conflict markers (`<<<<<<<, =====, >>>>>>>`). Verify all three governance files (Constitution + CLAUDE + AGENTS) cross-reference the new clause consistently.
4. **Commit + push to ALL upstreams.** Stage only the governance files (NEVER `git add -A` inside the constitution submodule — stray local artefacts MUST NOT enter governance). Commit with a message that cites the user mandate (verbatim quote) + the classification line per §11.4.17. Push to **every** configured remote of the constitution submodule. A commit that lives on one upstream but not others is a §11.4.26 violation equivalent to a §2.1 multi-upstream-push violation.
5. **Conflict resolution.** If `pull --ff-only` reports non-fast-forward, the merge MUST be performed carefully: inspect both sides, preserve the union of governance content (no clause silently dropped), re-classify per §11.4.17, validate per step 3. Force-push to “make conflicts go away” is FORBIDDEN

(§9.2). Nothing about the constitution may be broken, made faulty, corrupted, or rendered unusable by the merge.

6. **Post-merge validation + verification.** After the push lands, re-clone (or `git submodule update --remote --init`) in a throwaway worktree and re-run the consumer project's inheritance-cascade verifier (e.g. `scripts/verify-governance-cascade.sh`) to confirm the new clause reaches every owned submodule per CONST-047. Any cascade gap **MUST** be closed in the same change-window.
7. **Update the consuming project's pointer.** The consuming project's `.gitmodules`-tracked submodule pointer **MUST** be bumped to the new constitution HEAD in the SAME commit as any downstream cascade work; out-of-sync submodule pointers are a §11.4.26 violation.

Operational scope. This workflow applies regardless of who initiates the change (operator, primary agent, subagent per §11.4.20, automated lint suggestion). The workflow **CANNOT** be shortcut by “I’ll fetch later” or “I’ll push to the other upstreams in the next commit”. The constitution is the **single source of truth** for every project that imports it; allowing it to fragment across upstreams is the structural equivalent of a §11.4 PASS-bluff at the governance layer.

Cross-cutting reach. §11.4.26 composes with: §2 (single- entrypoint commit wrapper), §2.1 (multi-upstream push), §3 (submodule changes propagate through submodule commits first), §9.1 / §9.2 (data-safety + force-push authorization), §11.4.17 (universal-vs-project classification), §11.4.22 (lightweight doc-sync), §11.4.25 (full-automation coverage — the post-merge validation step is itself a form of automation coverage). It does NOT supersede them.

Classification: universal (per §11.4.17). No escape hatch. A constitution-submodule change that violates §11.4.26 is a release blocker for every consuming project, equivalent in severity to a force-push without §9.2 authorization.

§11.4.31 — Submodule-Dependency-Manifest Mandate (User mandate, 2026-05-15)

Forensic anchor — verbatim user mandate (2026-05-15):

“We **MUST HAVE** mechanism for each Submodule to determine / know what are its Submodule dependencies so new projects or palces we are incorporate them can add these Submodules to the project root and make them available! Suggested idea is configuration file with expected Submodules Git ssh urls perhaps? New project can read it, and recursively add each Submodule to the root of the project and install / expose it to everyone. This **MUST** be analyzed and applied. We **MUST** apply the best strategy for this which can be easily executed just by following our root Constitution, AGENTS.MD and CLAUDE.MD! Process this, extend out Constitution, AGENTS.MD and CLAUDE.MD with mandatory instructions and then process project root and all Submodules deep recursively so proper configuration Submodules dependency files are created! Document **EVERYTHING** and cover with all supported test types! Any kind of bluff is strictyl forbidden! All wotk **MUST** come with mechanism for validation and verification by creating proper proofs for all critical points!”

Tooling contract. A consuming project MUST be able to bootstrap its entire own-org dependency graph by:

1. Running `incorporate-submodule <ssh-url>` (canonical name; lives in the constitution submodule's `scripts/` or in a parent project's `bin` path).
2. The script:
 - Adds the supplied submodule at its declared canonical path (per CONST-051(C) flat vs grouped layout).
 - Reads the newly-added submodule's `helix-deps.yaml`.
 - For each declared dep, checks if it already exists at the consuming project's root; if missing, recurses (`incorporate-submodule` on the dep's `ssh_url`).
 - On conflict (same name declared at different ref by two submodules), aborts with a directed error pointing the operator at the conflicting declarations.
 - Emits a final manifest-of-manifests file at `<root>/.helix-manifest.yaml` listing every submodule + its declared deps, for audit + reproducibility.

Anti-bluff guarantee. Every manifest MUST be paired with a **verification proof**: a Challenge script (per §11.4.27 + CONST-050(B)) that:

- Bootstraps a throwaway consuming project from scratch in a temp directory.
- Runs `incorporate-submodule` against the manifest under test.
- Verifies the produced submodule layout matches the manifest's declarations (every dep present at its declared layout path; no extras; no missing).
- Runs the submodule's own test suite against the bootstrapped layout; asserts pass.
- Captures wire evidence (per §11.4.2) of every step.

A manifest without this verification proof is a §11.4.31 violation of equal severity to a §11.4 PASS-bluff at the dependency-graph layer.

Cascade requirement. Every owned-by-us submodule MUST ship `helix-deps.yaml` at its root, recursively (sub-submodules of submodules ship their own). The constitution submodule itself ships a manifest declaring its own deps (currently empty for the universal Constitution; project consumers may have project-specific extensions). The verifier (`scripts/verify-governance-cascade.sh` or its successor) MUST check every owned submodule for manifest presence.

Composition. §11.4.31 directly enables §11.4.28 / CONST-051(C) flat-layout enforcement: nested own-org submodule chains can be mechanically flattened because each submodule declares what it needs, and the incorporator places those deps at the root. Composes with §1 (manifest schema is itself tested for parse + validation), §3 (submodule-first commit discipline applies to manifest changes too), §11.4.12 (manifest changes regenerate any derived docs/diagrams), §11.4.17 (universal — no project-specific assumptions in the manifest format), §11.4.18 (manifest documented in script-doc + external user guide), §11.4.20 (subagent delegation for cross-submodule manifest authoring), §11.4.25 (manifest presence in coverage ledger), §11.4.26 (constitution-update workflow when extending manifest schema), §11.4.27 (manifest test matrix), §11.4.28 (this rule is its operational complement), §11.4.29 (manifests use `snake_case` names + canonical paths), §11.4.30 (manifest is tracked source — not a build artefact), CONST-047 (manifests cascade recursively).

Classification: universal (per §11.4.17). No escape hatch. Severity-equivalent to a §11.4 PASS-bluff at the dependency-graph layer.

§11.4.32 — Post-Constitution-Pull Validation Mandate (User mandate, 2026-05-15)

Forensic anchor — verbatim user mandate (2026-05-15):

“Every time we fetch and pull new changes on constitution Submodule we MUST process the whole project and all Submodule (deep recursively) for validation and verification taht every single rule or mandatory constraint is followed and respected! If it is not, IT MUST BE!”

Operative rule. Whenever a consuming project’s constitution submodule is fetched + pulled with **any** content change (rule addition, rule revision, gate addition, anchor reference, version bump), the consuming project MUST execute a full-project + recursive-submodule validation sweep BEFORE the new constitution HEAD is treated as canonical for any other work.

Validation sweep contract. The sweep is implemented as `scripts/verify-all-constitution-rules.sh` (canonical name) which:

1. Re-runs the existing governance-cascade verifier (`scripts/ verify-governance-cascade.sh`) covering every §11.9 + CONST-* anchor across every owned submodule (recursive per CONST-047).
2. For each rule whose enforcement gate is implementable programmatically (e.g., CONST-053 `.gitignore`-pattern audit; CONST-051(C) nested-own-org-chain audit; CONST-052 case-conformance audit; CONST-050(A) mock-path-from-production- code audit; CONST-035 anti-bluff smoke-scan), the sweep runs the corresponding gate against the post-pull tree.
3. Any failure produces a directed FAIL entry naming the rule (e.g., FAIL: CONST-053 – *.log tracked at HelixCode/foo.log)
 - the canonical fix.
4. Failures populate the project’s Issues tracker per §11.4.15 with Status: Reopened (since a previously-passing audit now fails) and Type: Bug (since real codebase state violates the constitution).
5. The agent or operator MUST resolve every FAIL before treating the new constitution HEAD as canonical — closure of each reopened item per the existing §11.4 anti-bluff covenant (positive-evidence-only, captured wire evidence).

Pull-time invocation. `git submodule update --remote constitution` MUST trigger the sweep automatically (post-update hook OR the operator’s commit wrapper invokes it as part of any commit that advances the constitution submodule pointer). Operator-explicit manual invocation MUST also be available (`./scripts/verify-all-constitution-rules.sh`).

Anti-bluff guarantee. A sweep that exits PASS without actually running every implementable gate is a §11.4.32 violation. The sweep’s own meta-test (paired mutation, §1.1) MUST plant a known violation of each enforced gate and assert the sweep reports FAIL for the planted gate.

Composition. §11.4.32 is the **enforcement engine** for every other §11.4.x and CONST-NNN rule. Without it, new rules cascade as anchors but never get enforced in the codebase. Composes with every rule that has a programmatic gate: §1, §1.1, §11.4.10, §11.4.12, §11.4.15, §11.4.16, §11.4.18, §11.4.20, §11.4.22, §11.4.24, §11.4.25, §11.4.26, §11.4.27, §11.4.28, §11.4.29, §11.4.30, §11.4.31, CONST-035,

CONST-038, CONST-042, CONST-043, CONST-044, CONST-045, CONST-046, CONST-047, CONST-048, CONST-049, CONST-050, CONST-051, CONST-052, CONST-053, CONST-054.

Classification: universal (per §11.4.17). No escape hatch. Severity-equivalent to a §11.4 PASS-bluff at the constitutional- enforcement layer — without §11.4.32, every other rule is a decorative anchor rather than an enforced gate.

§11.4.30 — .gitignore + No-Versioned-Build-Artifacts Mandate (User mandate, 2026-05-15)

Forensic anchor — verbatim user mandate (2026-05-15):

“every project module, every Submodule, every service and application MUST HAVE proper .gitignore file! We MUST NOT git version build artifacts, cache files, tmp files, main .env file(s) or any files containing sensitive data, API keys or token! Any build derivative which we can recreate by executing proper mechanism for generating MUST NOT be versioned! We MUST pay attention what is going to be committed every time we are preparing to execute commit! If any violation is detected it MUST be fixed before commit is executed!”

Operative rule. Every project module, owned-by-us submodule, service, and application under this Constitution MUST ship a proper .gitignore covering its full set of forbidden-from-version- control patterns. The following file/directory classes MUST NEVER appear in version control:

1. **Build artefacts** — anything produced by a build / compile / package step that can be regenerated from sources:
 - Binaries (/bin/, /build/, /dist/, /out/, target/, *.exe, *.dll, *.so, *.dylib, *.a, *.o, *.class, *.pyc).
 - Generated source files where the generator is committed (e.g. protobuf .pb.go may be checked in OR ignored — project decides — but never both).
 - Bundled assets where the bundler is committed.
2. **Cache files** — anything a tool regenerates on demand:
 - __pycache__/, .pytest_cache/, .mypy_cache/, .ruff_cache/, node_modules/, .next/, .nuxt/, .cache/, .gradle/, .idea/cache/, .vscode-test/, target/, .terraform/, language-server caches.
3. **Temp files** — *.tmp, *.swp, *.swo, *~, .DS_Store, Thumbs.db, IDE backup files, *.orig, *.rej.
4. **Sensitive-data files** — anything containing credentials, tokens, keys, or personal data:
 - .env, .env.* (allow .env.example / .env.sample / .env.template with placeholder values only — never real secrets even as examples).
 - *.pem, *.key, *.crt, *.p12, *.pfx, id_rsa*, id_ed25519*, *.kdbx, .netrc, secrets/ directory trees.

- `api_keys.sh`, any file containing `BEGIN PRIVATE KEY`, credential tokens, or session cookies.

5. **Generated reports / logs** — `*.log`, `coverage.out`, `*.coverage`, `htmlcov/`, `*.gcda`, `*.gcno`, screenshots/ recordings that aren't reference assets.
6. **OS / IDE / personal-state files** — `.idea/`, `.vscode/` (except `shareable` `.vscode/settings.json` etc. if explicitly project-shared), `.history/`, `.svn/`, editor-state files.
7. **Test playback artefacts** (§11.4.14) — every test's runtime capture goes under an ignored evidence directory unless the capture IS the reference asset.

Anti-bluff invariant. Per CONST-035 / §11.4 the presence of a `.gitignore` line alone is not sufficient — the project MUST also verify that no file matching the forbidden patterns is currently tracked. A `.gitignore` that lists `*.log` while `app.log` sits tracked in the repo is a §11.4.30 violation of equal severity to no `.gitignore` at all.

Pre-commit attention. Every commit author (human OR agent) MUST inspect `git diff --staged` and `git status` BEFORE executing the commit. If a staged path matches any forbidden class, the commit MUST be aborted and the issue fixed (un-stage the path, add to `.gitignore`, scrub if already-tracked). Gate CM-GITIGNORE-PRECOMMIT-AUDIT: pre-commit hook (or commit wrapper in §2) inspects every staged path against the forbidden-pattern matrix; hits abort the commit with a directed error pointing the operator at the specific offending path and the canonical fix. Paired mutation (§1.1): stage a `*.log` → gate FAILs.

Cascade reach. This rule applies recursively to every owned-by- us submodule per CONST-047 and to every owned-submodule's nested non-owned trees within their working copy. The `.gitignore` files themselves are project-specific in their concrete patterns but universally bound to the rule's normative force.

Secret-leak intersection. §11.4.30 composes tightly with §11.4.10 (Credentials-handling mandate / CONST-042) and §12.1. A `.env` leak is BOTH a §11.4.30 violation (build/sensitive file versioned) AND a §11.4.10 violation (credential leak). The combined severity is **release-blocker requiring rotation + post-mortem** per §11.4.10.

Coverage of “recreatable” content. When in doubt about whether something is a build derivative: if there exists a documented, scripted, or automated mechanism that recreates the file from sources, the file is a build derivative and MUST be ignored. The committed sources MUST include the generator (Makefile target, npm script, codegen invocation), so any downstream consumer regenerates the artefact on demand.

Classification: universal (per §11.4.17). No escape hatch beyond the explicit exceptions enumerated above (e.g., `.env.example` placeholder files). Severity-equivalent to a §11.4 PASS-bluff at the repository-hygiene layer — a tracked build artefact silently drifts vs. its source over time, producing the same class of “works-on-my-machine” surprise that the §11.4 anti-bluff covenant forbids at the test layer.

Composition. §11.4.30 composes with §1 (test coverage — ignored files don't get spuriously included in coverage stats), §2 (commit-wrapper enforces the pre-commit gate), §9.1 (destructive-operation safeguards — never `git clean` an operator's

working tree to “fix” a violation), §11.4.10 (credentials-handling), §11.4.12 (auto-generated docs sync — generators committed, outputs ignored only if recreatable on demand from the same generator), §11.4.17 (universal — applies to every consuming project), §11.4.18 (script docs — generator scripts documented), §11.4.20 (subagent delegation for cross-submodule .gitignore audit sweeps), §11.4.25 (coverage ledger lists .gitignore presence per submodule), §11.4.26 (constitution-update workflow — the constitution submodule’s own .gitignore complies), §11.4.27 (no-fakes: test fixtures that ARE the reference asset are tracked; runtime captures are ignored), §11.4.28 (submodules-as-equal-codebase: each submodule’s .gitignore audited on equal basis), §11.4.29 (the renamed paths in the lowercase migration MUST also be properly ignored where applicable), CONST-047 (recursive governance reach).

§11.4.29 — Lowercase-Snake_Case-Naming Mandate (User mandate, 2026-05-15)

Forensic anchor — verbatim user mandate (2026-05-15):

“naming convention for Submodules and directories (applied deep into hierarchy recursively) - all directories and Submodules MSUT HAVE lowercase names with space separator between the words of ‘_’ character (snake-case)! All existing Submodules and directories which are not following this rule MUST BE renamed! However, since this will most likely break some of the functionalities renaming we do MUST BE applied to all references to particular Submodule or directory! Everywhere where particular Submodule directory are referenced proper updates MUST BE applied - all configuration files, documentation and relevant materials, links to Submodules and directories, source code that points to them, etc. There MUST BE reasonable exceptions for this rules - source code for programming languages or Submodules which apply different naming convention - Android, Java, Kotlin and others. Root directory for such applications, services or Submdoules can follow OUR convention, but EVERYTHING inside still MUST follow language / technology specific rules! We apply this rules per common sense basis and it MUST NOT be the cause of bigger issues such as technology breaking! Upstreams directory which all of our projects and Submodules have MUST BE renamed to the lowercase letters too, however root project containing the install_upstreams system command (it is exported in out paths in our .bashrc or .zshrc) MUST BE updated to fully work with both Upstreams and upstreams directory. That change if it is not already applied MUST BE done, committed and pushed! ... NOTE: Rules lowercase / snake-case do apply to all project files as well and references to it and from them! Every change done MUST BE covered with all supported test types, full automation tests for validation and verification and fully applied and followed anti-bluff policy and all anti-bluff rules!”

Operative rule. Every directory, submodule, and file under the parent project’s working tree MUST use a **lowercase, snake_case** name (ASCII letters / digits / underscores, words separated by _). Existing names that violate the rule (HelixCode/, Challenges/, Containers/, HelixAgent/, HelixQA/, Security/, Github-Pages-Website/, Upstreams/, Dependencies/, etc.) MUST be renamed as part of the migration window opened by this clause. Every reference in the codebase MUST be updated atomically with the rename: configuration files,

documentation, user manuals, diagrams, scripts, source-code imports, links, governance files. **Reference drift after a rename is a §11.4.29 violation** of equal severity to the rename itself.

Exceptions (common-sense scope). The rule MUST NOT break language-/technology-specific conventions:

- **Programming-language source roots** that mandate a specific case (Java / Kotlin package paths, Android resource folders, Apple framework directories, C# / Swift project layouts) keep their language-mandated names. The submodule's root directory follows our convention; the language-specific subtree inside follows its own.
- **Vendor / upstream submodules** (third-party orgs not in our owned set) keep their upstream-mandated names — we MUST NOT rename a third party's repo.
- **Build-tooling artefacts** (node_modules/, __pycache__/, .git/, target/, build/, bin/) keep their tool-mandated names.

When in doubt, the test “does renaming break the technology?” trumps the snake_case rule. The §11.4.29 spirit is operator- ergonomics + reference-discoverability, not pedantic uniformity at the cost of technology compatibility.

Upstreams/ → upstreams/ transition. The constitution submodule's installer (install_upstreams.sh) — exported on operator paths via .bashrc / .zshrc — MUST support **both** Upstreams/ and upstreams/ directory layouts during the migration window, so existing checkouts keep working while consuming projects rename at their own pace. The installer reads whichever directory exists; if both exist the lowercase wins. After every project under this Constitution has migrated, the uppercase fallback MAY be retired by a deliberate amendment, but the migration window remains open as long as ANY owned project still ships the uppercase form.

Project-Toolkit Upstreamable submodule synchronisation. The Upstreamable / Project-Toolkit machinery that propagates governance into every consuming project MUST be fetched + pulled before any rename batch, and MUST itself comply with this rule. Any Upstreamable submodule lacking BOTH-directory support is a release blocker for the rename program.

Test coverage of renames. Every batch of renames MUST ship with: (i) a regression test that verifies every reference to the renamed entity now resolves to the new name (no stale references left); (ii) a full CONST-050(B) test-type matrix run against the post-rename tree; (iii) anti-bluff (CONST-035) wire-evidence captured during the runtime verification. A rename batch without all three is a §11.4.29 violation.

Cascade reach. This rule applies recursively through every owned-by-us submodule layer (per CONST-047) and applies equally to the consuming project and its owned submodules (per CONST-051). Per CONST-051(C) — dependencies at the parent root — the renamed paths MUST be the only canonical location; old-name aliases remain only as transitional symlinks if absolutely required, and those symlinks MUST be removed on next-N-cycle review (project- configurable, recommended ≥3 release cycles).

Phased execution. Because the rename touches every reference, the execution MUST be planned as fine-grained phases per the operator's explicit instruction: comprehensive brainstorming, phase-divided plan, fine-grained tasks/subtasks with

enormous detail, every change covered by every applicable test type. The phases run in parallel with mainstream work (§11.4.20 subagent delegation is the natural fit for cross-submodule rename sweeps).

Classification: universal (per §11.4.17). No escape hatch beyond the explicit common-sense exceptions enumerated above. Severity-equivalent to a §11.4 PASS-bluff at the reference-integrity layer — a half-completed rename that leaves broken references is worse than no rename at all because it silently breaks consumers.

Composition. §11.4.29 composes with §1 (four-layer floor for every rename batch), §1.1 (paired mutation: rename without updating a reference → gate FAILs), §11.4.12 (regenerate auto-generated docs on rename), §11.4.17 (universal — no project-specific assumptions), §11.4.18 (script-doc-sync after renamed script paths), §11.4.20 (subagent delegation for cross-cutting rename sweeps), §11.4.25 (every renamed path appears in coverage ledger), §11.4.26 (constitution-submodule rename pipeline), §11.4.27 (rename-touch test types apply across the matrix), §11.4.28 (owned submodule renames cascade per CONST-047), CONST-047 (recursive governance reach).

§11.4.28 — Submodules-As-Equal-Codebase + Decoupling + Dependency-Layout Mandate (User mandate, 2026-05-15)

Forensic anchor — verbatim user mandate (2026-05-15):

“All existing Submodules in the project that we are controlling and belong to some our organizations (vasic-digital, HelixDevelopment, red-elf, ATMOSphere1234321, Bear-Suite, BoatOS123456, Helix-Flow, Helix-Track, Server-Factory — we can ALWAYS check dynamically using GitHub and GitLab CLIs) are equal parts of the project’s codebase! We MUST work on that code as much as we do with main project’s codebase! All on equal basis! Equally important! We MUST take it into the account, analyze it, extend it, create missing tests, do full testing of it, fill the gaps (if any), fix any issues that we discover or they pop-up, write and extend the documentation, user guides, manulas, diagrams, graphs, SQL definitions, Website(s) and all other relevant materials! We MUST NEVER modify Submodules to bring into them any project specific context since they all MUST BE ALWAYS fully decoupled, project not-aware, fully reusable and modular (by any other project(s)), completely testable! All Submodule dependencies that are used by Submodule MUST BE accessed from the root of the project! We MUST NOT have nested Submodule dependencies but accessing each from proper location from the root of the project — directly from project’s root project_name/ submodule_name or some more proper structure project_name/ submodules/submodule_name! This MUST BE heavily enforced and respected so no chaos and mess is created with various dependencies we may have!”

Operative rule. Three cooperating invariants govern every consuming project’s relationship with its owned-by-us submodules (those whose upstream origin lives under one of the operator-listed orgs — vasic-digital, HelixDevelopment, red-elf, ATMOSphere1234321, Bear-Suite, BoatOS123456, Helix-Flow, Helix-Track, Server-Factory — or any additional org the operator subsequently authorises, the canonical list discoverable at any time via `gh org list / glab / the orgs’ public APIs`):

(A) Equal-codebase invariant. Every owned-by-us submodule is an **equal part** of the consuming project's codebase. The consuming project's engineering practice — analysis, extension, test creation, gap-filling, bug-fix, documentation (user manual, guides, diagrams, graphs, SQL definitions, website pages, any other authoring surface) — applies to each owned submodule on equal basis. A round of work that improves only the main project while leaving an owned-submodule deficiency unaddressed is a §11.4.28 violation, severity-equivalent to a §11.4 PASS-bluff at the project-scope layer. Coverage ledgers (§11.4.25) **MUST** list every owned submodule as an in-scope target. CONST-047 (recursive cascade) governs the propagation of governance changes; §11.4.28 is the engineering-content counterpart that mandates the same attention to non-governance content.

(B) Decoupling / reusability invariant. Owned submodules **MUST** remain **fully decoupled** from any specific consuming project. No project-specific context, hardcoded paths, hostnames, asset names, naming schemes, or runtime assumptions may be introduced into an owned submodule's source tree. Every owned submodule **MUST** be:

- **Project-not-aware** — its code, tests, and docs make no reference to which parent project consumes it.
- **Fully reusable** — any future Helix-family or third-party project must be able to consume the submodule unmodified.
- **Modular** — its public surface is the only documented integration contract; internal layout may evolve without breaking consumers.
- **Completely testable** — every public surface has standalone tests (per the §11.4.27 100%-test-type matrix) that pass when the submodule is checked out as a standalone repo, without any parent-project rigging.

A commit that adds `project_name/...` strings, hostnames belonging to a specific deployment, or other parent-project context to a submodule's source tree is a §11.4.28 violation. The honest path when a submodule needs information from the parent project is configuration injection (env var, config file, constructor parameter) — never a hardcoded reach into the parent's tree.

(C) Dependency-layout invariant. Every dependency that an owned submodule itself consumes **MUST** be accessible **from the root of the parent project** at one of two canonical paths:

```
<project_root>/<submodule_name>/          # flat layout
<project_root>/submodules/<submodule_name>/ # grouped layout
```

Nested-submodule chains are FORBIDDEN. A submodule **MUST NOT** have its own `.gitmodules` entries that pull in further owned- by-us repos (transitive own-org submodule recursion). Every dependency required by submodule X **MUST** be added to the parent project at the canonical path above; X reaches it via documented import / SDK path / runtime resolver — never via its own nested submodule pointer.

Rationale: nested own-org submodule chains cause version-drift chaos (two consumers of `LLMsVerifier` end up at different SHAs because each parent picked a different transitive path). The flat / grouped layout makes the consuming project's submodule graph a tree-of-depth-1, which any developer can audit at a glance via `git submodule status` from the project root.

Third-party submodules (not under our orgs) are exempt — they MAY appear at any depth as the upstream’s structure dictates. The invariant applies only to our owned set.

Audit + enforcement.

- Gate CM-OWNED-SUBMODULE-EQUAL-ENGINEERING (project-side): every release-gate sweep verifies each owned submodule has current test runs, coverage entries (§11.4.25), and documentation freshness on par with the main project. Stale submodules surface as §11.4.28 violations (Status: Operator-blocked or In progress per §11.4.21 — never “ignored”).
- Gate CM-OWNED-SUBMODULE-DECOUPLING (submodule-side): every owned submodule’s pre-commit hook (or equivalent) greps the staged diff for parent-project names / hostnames / asset names. Hits abort the commit until refactored to configuration injection.
- Gate CM-OWNED-SUBMODULE-LAYOUT (project-side): the parent project’s pre-merge sweep verifies (i) every owned submodule sits at <root>/<name>/ or <root>/submodules/<name>/,
 1. no owned submodule contains a nested .gitmodules entry whose upstream is in our org list, and (iii) every dependency declared by an owned submodule has a corresponding parent-project submodule entry at the canonical path.
- Paired mutations (§1.1) for all three gates: plant the forbidden pattern → gate FAILs; restore → gate PASSes.

Workflow integration. Honoring §11.4.28 in practice:

- Engineering rounds plan in submodule-aware tranches: “improve feature X in main; in same round, audit + improve the X- related surface in submodule Y; cascade governance per CONST-047 as usual”.
- The §11.4.25 coverage ledger row format is extended with a submodule column so coverage rolls up across the whole owned set.
- Cross-submodule refactors that touch shared types or interfaces ship as a single change-window: parent + every affected owned submodule advanced in lockstep (§3 submodule- first-commit discipline + §2.1 multi-upstream push).

Classification: universal (per §11.4.17). No escape hatch. Composes with: §1 (four-layer test floor reaches submodules too), §1.1 (false-positive immunity), §3 (submodule changes propagate through submodule commits first), §11.4.17 (universal-vs-project classification), §11.4.20 (subagent delegation for the cross-submodule audit sweep), §11.4.25 (full-automation coverage ledger expanded to submodules), §11.4.26 (constitution-update workflow’s submodule-pointer-bump step), §11.4.27 (100%-test- type coverage applies to every submodule’s standalone surface), CONST-047 (recursive governance cascade). A round of work that violates §11.4.28 is a release blocker for the consuming project, severity-equivalent to a §11.4 PASS-bluff at the codebase-completeness layer.

§11.4.27 — No-Fakes-Beyond-Unit-Tests + 100%-Test-Type-Coverage Mandate (User mandate, 2026-05-15)

Forensic anchor — verbatim user mandate (2026-05-15):

“Mocks, stubs, placeholders, TODOs or FIXMEs are allowed to exist ONLY in Unit tests! All other test types MUST interact with real fully implemented System! No fakes, empty implementations or bluffing is allowed of any kind! All codebase of the project MUST BE 100% covered with every supported test type: unit tests, integration tests, e2e tests, full automation tests, security tests, ddos tests, scaling tests, chaos tests, stress tests, performance tests, benchmarking tests, ui tests, ux tests, Challenges (fully incorporating our Challenges Submodule — <https://github.com/vasic-digital/Challenges>). EVERYTHING MUST BE tested using HelixQA (fully incorporating HelixQA Submodule — <https://github.com/HelixDevelopment/HelixQA>). HelixQA MUST BE used with all possible written tests suites (test banks) for every applications, service, platform, etc and execution of the full HelixQA QA autonomous sessions! All required dependency Submodules MUST BE added into the project as well (fully recursive!!!).”

Operative rule. Two cooperating invariants:

(A) No-fakes-beyond-unit-tests. Mocks, stubs, fakes, placeholders, in-memory facsimile implementations, TODO, FIXME, “for now”, “in production this would”, or any empty-implementation pattern are PERMITTED only inside unit-test sources (e.g., *_test.go files invoked WITHOUT the integration build tag; tests/unit/; equivalent per-language conventions). Every other test type — integration, end-to-end, full automation, security, DDoS, scaling, chaos, stress, performance, benchmarking, UI, UX, Challenges, HelixQA suites — MUST exercise the **real, fully implemented system** against real infrastructure (real databases, real HTTP endpoints, real containers, real downstream services, real captured devices). A non-unit test that imports mocks, in-memory repositories, fabricated provider responses, or placeholder structs is a §11.4.27 violation regardless of how green its summary line looks — severity-equivalent to a §11.4 PASS-bluff. The same prohibition extends to **production code**: no mock import path may be reachable from any code path that runs in production binaries. Pre-build gate CM-NO-FAKES-BEYOND-UNIT-TESTS scans the non-unit test trees and production trees for the forbidden patterns; paired mutation plants a fake → gate FAILs.

(B) 100% test-type coverage with every supported type. Every project under this Constitution MUST cover its entire codebase with **every supported test type** the project’s domain warrants:

1. **Unit** — fast, isolated, mocks permitted per (A).
2. **Integration** — multi-component, no mocks, real backing services.
3. **End-to-end (E2E)** — full user-flow exercise on target topology.
4. **Full automation** — orchestrated suites exercising every feature × platform combination (§11.4.25 coverage ledger).
5. **Security** — authn/authz boundaries, secret-leak scans (§11.4.10), input-fuzzing, dependency-CVE scanning, threat- model verification.
6. **DDoS** — request-flood resilience at advertised throughput tier, with rate-limit + back-pressure assertions.
7. **Scaling** — horizontal + vertical scale behaviour under linear load growth, including replica add/remove transitions.
8. **Chaos** — controlled failure injection (network partition, process kill, disk full, clock skew) verifying graceful degradation paths.
9. **Stress** — sustained load above advertised tier, asserting bounded resource exhaustion + clean recovery.
10. **Performance** — latency / throughput / tail-latency invariants vs SLO baselines.

11. **Benchmarking** — micro + macro benchmark suites with historical p95-drift detection (§11.4.24 build-resource composition).
12. **UI** — visual-regression + DOM-state + interaction-flow coverage on every target platform's UI surface.
13. **UX** — flow-correctness + accessibility + i18n + visual- cue ordering (§11.4.23 composition).
14. **Challenges** — vasic-digital/Challenges submodule fully incorporated; per-feature Challenge scripts covering real user use-cases with captured runtime evidence.
15. **HelixQA** — HelixDevelopment/HelixQA submodule fully incorporated; ALL written test banks executed; full autonomous QA sessions run as part of release gates.

Per-type 100% coverage means: for every feature × platform cell in the §11.4.25 coverage ledger, the cell carries a verified PASS evidence pointer for **each supported test type the cell warrants** (a CLI-only feature warrants unit + integration + E2E + full-automation + security + chaos + stress + performance + Challenges + HelixQA; a UI feature additionally warrants UI + UX; a network service additionally warrants DDoS + scaling + benchmarking). Gaps are tracked per §11.4.15 with explicit UNCONFIRMED: / PENDING_FORENSICS: / OPERATOR-BLOCKED: reasons.

Required submodule incorporation (recursive). Every project consuming this Constitution MUST add the following as Git submodules (or vendored equivalents authorised by the operator), fully recursively per CONST-047 / §11.4.x cascade:

- Challenges — git@github.com:vasic-digital/Challenges.git
- HelixQA — git@github.com:HelixDevelopment/HelixQA.git
- Any additional functionality submodules under vasic-digital/* or HelixDevelopment/* orgs that the consuming project depends on (do not duplicate work the organisations already maintain).

Submodule pointers MUST be bumped to upstream HEAD in the SAME commit as any dependent cascade work (§11.4.26 step 7). Pointer drift = §11.4.27 violation.

HelixQA autonomous sessions. “Full HelixQA QA autonomous sessions” means: HelixQA's orchestrator drives the consuming project's running surface end-to-end, executing every test bank the project has registered with it, capturing wire evidence per check (status + body bytes + body-head + duration + screenshots for UI), and producing a session report that is itself committed via the §11.4.22 lightweight doc-sync wrapper. A green autonomous-session report without captured wire evidence is a §11.4 PASS-bluff at the QA-orchestration layer.

Composition. §11.4.27 composes with: §1 (four-layer floor), §7.1 (positive-evidence-only), §11.4.1 (FAIL-bluffs forbidden), §11.4.2 (recorded-evidence), §11.4.3 (per-topology dispatch), §11.4.6 (no-guessing), §11.4.10 (credentials-handling — security test category), §11.4.15 / §11.4.16 (status + type tracking), §11.4.17 (universal-vs-project classification), §11.4.20 (subagent delegation for orchestrating the test-type matrix), §11.4.22 (lightweight doc-sync for coverage ledgers and session reports), §11.4.25 (full-automation-coverage — §11.4.27 is its strict expansion into per-type-of-test territory). It does NOT supersede them. CONST-047 (recursive submodule application) governs the cascade reach.

Classification: universal (per §11.4.17). No escape hatch. A project shipping with mocks reachable from non-unit-test paths OR missing required test-type coverage is **not delivering a stable product** — release blocker for every consuming project, severity-equivalent to a §11.4 PASS-bluff at the release-gate layer.

§11.4.33 — Type-aware closure-status vocabulary (User mandate, 2026-05-15)

Forensic anchor — direct user mandate (verbatim, 2026-05-15):

“make sure we use proper wording for the workable item completion status in Issues, Issues_Summary and Fixed docs: - Fixed -> into the proper wording depending on the workable item type (task, feature, and so on) - for example: Fixed, Implemented, Completed. Add this important detail in our root Constitution, CLAUDE.MD and AGENTS.MD so it is ALWAYS respected and followed!”

Classification: §11.4.17-classified **universal** — naming discipline applies to every project that tracks work items by type. Bug closures aren’t “implemented”, feature deliveries aren’t “fixed”, and infrastructure refactors aren’t either. Type-mismatched closure vocabulary is a quiet but persistent semantic-drift bluff at the documentation layer — it makes greppable releases lie about what shipped.

The mandate. §11.4.15 defined the lifecycle Status closed-set including the closure terminal value Fixed (→ Fixed.md). §11.4.16 defined the Type closed-set {Bug | Feature | Task}. §11.4.33 binds the two: the closure terminal value MUST agree with the item Type, drawn from this 3-element closed map:

Item **Type:**	Closure **Status:** value
Bug	Fixed (→ Fixed.md)
Feature	Implemented (→ Fixed.md)
Task	Completed (→ Fixed.md)

The (→ Fixed.md) suffix is preserved across all three so the existing migration-discipline tooling (Issues.md → Fixed.md atomic move per §11.4.19) continues to work without per-Type branching. Generators (generate_issues_summary.sh, generate_fixed_summary.sh, status-counter helpers, the §11.4.23 colorizer) MUST treat the three terminal values as semantically equivalent (all map to “closed, positive evidence captured”) while preserving the literal in the emitted document.

No escape hatch. Closing a Feature with Fixed (→ Fixed.md) or a Task with Implemented (→ Fixed.md) is a §11.4.33 violation. Pre-build gate (recommended, per consuming project) CM-CLOSURE- VOCAB-TYPE-AWARE walks every Fixed.md heading + every Issues.md heading whose ****Status:**** is one of the three terminal values, and asserts the Status value matches the item’s ****Type:**** per the table. Paired mutation flips a Bug entry’s status from Fixed (→ Fixed.md) to Implemented (→ Fixed.md) → gate FAILs.

Propagation. Composes with §11.4.15 (item-status tracking), §11.4.16 (item-type tracking), §11.4.19 (Fixed-document column alignment), §11.4.23 (colorisation — the closed-state palette applies regardless of which of the three terminal words is used).

§11.4.34 — Reopened-source attribution mandate (User mandate, 2026-05-15)

Forensic anchor — direct user mandate (verbatim, 2026-05-15):

“when we reopen some workable item (bug, task or feature) we MUST HAVE details on who did reopened this and why (- Reopened status needs: by who, AI or User + details in main Issues doc.). For example, reopened by AI or reopened by real User, why - test failed, manual testing detected problem, etc. Adapt our docs for this - Issues, Issues_Summary and Fixed.”

Classification: §11.4.17-classified **universal** — every project that uses §11.4.15’s Reopened lifecycle value benefits from knowing the reopen source + cause. Without this, reopen-thrash patterns (the same item bouncing between Fixed and Reopened across cycles) cannot be diagnosed, and the §11.4.7 demotion-evidence rule loses its forensic counterparty (you cannot evaluate whether a reopen was operator-side observation or agent-side over-restoration without provenance).

The mandate. Every Issues.md (or equivalent project tracker) heading whose ****Status:**** value is Reopened MUST carry, within 8 non-blank lines of the heading, a ****Reopened-Details:**** line that captures four sub-facts:

- **By:** AI or User (the source-of-truth observer who flipped the status). AI covers in-loop reopens (test failure, gate regression, captured-evidence retrospect). User covers operator- side observations (manual testing, end-user report, design reconsideration).
- **On:** ISO date (YYYY-MM-DD).
- **Reason:** one-line cause classification — chosen from a closed vocabulary { test-failed | manual-testing-detected | captured-evidence-contradicts | end-user-report | cycle-re-discovered | design-reconsidered }. Other values are permitted with explicit Reason: <free text> annotation but the closed list MUST be tried first.
- **Evidence:** path to or short description of the captured artefact that justifies the reopen — log file, recording, gate failure ID, operator quote, etc. Reopens without evidence are §11.4.6 / §11.4.7 violations: the reopen IS a demotion-from-Fixed classification change, and demotion requires positive evidence captured under the conditions that re-exposed the defect.

The Issues_Summary.md (or equivalent) Status column MUST distinguish the four Reopened sub-states by source so a sweep query for “reopens by AI in the last 30 days” is mechanically possible. Suggested column rendering: Reopened (AI: test-failed) vs Reopened (User: manual-testing).

No escape hatch. A Reopened entry without ****Reopened-Details:**** is a §11.4.34 violation. Pre-build gate (recommended, per consuming project) CM-ITEM-REOPENED-DETAILS mirrors CM-ITEM-OPERATOR-BLOCKED-DETAILS (Phase 39.AT pattern): walks every actionable heading, scans 8 non-blank lines for a ****Status:**** Reopened line, then continues scanning for ****Reopened-Details:****. Missing details line emits WARN initially, hardens to FAIL once backlog is fully populated.

Propagation. Composes with §11.4.6 (no-guessing — the Reason must be drawn from the closed vocabulary or explicitly annotated; no likely/probably/maybe causes), §11.4.7 (demotion-evidence — reopen IS a demotion from Fixed), §11.4.15 (item-status tracking — extends the Reopened value’s discipline), §11.4.21 (Operator- blocked discipline — same audit-line pattern with the same gate shape).

§11.4.35 — Canonical-root inheritance clarity (User mandate, 2026-05-15)

Forensic anchor — direct user mandate (verbatim, 2026-05-15):

“Parent AGENTS.MD or CLAUDE.MD are located under constitution directory (Submodule) containing these parent (root) which files we are inheriting inside the project! Pay attention to this and make sure you ALWAYS follow this rule! Same applies to the root Constitution file!”

“If needed (and we think it is needed it seems) add into the root Constitution, AGENTS.MD and CLAUDE.MD located in constitution directory (Submodule) which we are inheriting that the root (parent) Constitution, AGENTS.MD and CLAUDE.MD are not the ones in root of the project but inside the constitution directory (Submodule). We are inheriting it and if project specific rules have to be added or project specific constraints which are not universal and reusable, then they go into the Constitution, CLAUDE.MD and AGENTS.MD of the project directory itself!”

Classification: §11.4.17-classified **universal** — every project that consumes this constitution as a Git submodule **MUST** unambiguously distinguish the **canonical root** (the constitution submodule’s three files) from the **consumer extensions** (the project’s own three files).

The defect this anchor closes. Loose terminology — “parent CLAUDE.md”, “root CLAUDE.md”, “main constitution” — used without a file-path anchor causes a quiet but persistent confusion between the inheriting copy and the source-of-truth. AI agents have already been observed editing the consumer-side CLAUDE.md when the User’s intent was to extend the canonical layer (or vice-versa), causing universal rules to leak into project-specific copies and project-specific rules to be falsely promoted as universal. The defect compounds: once a rule is misfiled, future propagation gates treat the misfile as canonical and silently spread it.

The mandate. The three files in **this constitution submodule** (constitution/Constitution.md, constitution/CLAUDE.md, constitution/AGENTS.md) are the **canonical root** — also called the **parent** files. They contain only universal rules per §11.4.17 — rules that any project consuming this submodule benefits from.

The three files in **the consuming project’s repository root** (<project-root>/CLAUDE.md, <project-root>/AGENTS.md, optionally <project-root>/Constitution.md or equivalent) are the **consumer extensions**. They **MUST** start with the inheritance pointer per the existing constitution/CLAUDE.md “How inheritance works” section. They contain only project-specific rules per §11.4.17 — rules that reference particular hardware, vendor names, regulatory regions, internal asset names, or project-private conventions.

Operative invariants that follow:

1. **When in doubt about which file to edit:** if the rule is reusable across any project, edit the constitution submodule's file. If the rule references project-private specifics (a hardware revision, a vendor SDK version, a regional constraint, a particular service), edit the consumer's file. Default to consumer-side when uncertain (per §11.4.17 — narrower scope is cheap to widen later, the reverse is expensive).
2. **Terminology anchor.** When prose in any file in this constitutional family references “the parent CLAUDE.md” or “the root Constitution,” the referent is the constitution-submodule file at `constitution/<filename>`, never the consumer's file. When it references “the project CLAUDE.md” or “this project's AGENTS.md,” the referent is the consumer-side file at `<project-root>/<filename>`. AI agents reading this constitution **MUST** resolve ambiguous pronouns (“the CLAUDE.md”, “the constitution”) via this rule.
3. **Universal rules added to the consumer side are misfiled.** Per §11.4.17 commit-time check, before adding a **MUST** to the consumer file, the author assesses whether it would benefit other projects that consume this constitution. If yes, lift it to the constitution submodule first.
4. **Project-specific rules added to the constitution submodule are misfiled.** Per §11.4.17, before adding a **MUST** to the constitution submodule, the author assesses whether it references project-private specifics. If yes, demote it to the consumer file (or genericise it to a universal form first).
5. **Propagation gates target both layers but their contents differ.** The existing `CM-COVENANT-114-N-PROPAGATION` gate family verifies that anchor `TEXT` for §11.4.N exists in **EVERY** `CLAUDE.md` and `AGENTS.md` across the project. The constitution-submodule files carry the **canonical** text; the consumer files carry **compact propagation pointers** that reference the canonical authority by file path. Both forms count as valid presence; the gate is satisfied when any form is found.
6. **@constitution/CLAUDE.md import (Claude-Code-style) and the pointer-block fallback (Aider/Codex/Gemini-style) are equivalent.** The consumer's `CLAUDE.md` **MUST** start with one of:
 - The native import: `@constitution/CLAUDE.md`
 - The portable pointer-block per the `## INHERITED FROM constitution/CLAUDE.md` heading defined in `constitution/CLAUDE.md` “How inheritance works”.
7. **No silent demotion or silent promotion.** Moving a rule between layers **MUST** be a visible commit — `git mv` of a section if it's a clean clone, or an explicit “Lifted from to constitution per §11.4.35” / “Demoted from constitution to per §11.4.35” line in the commit message. AI agents **MUST NOT** silently re-author a §11.4.X anchor in the wrong layer and call it propagation.

Pre-build gate (recommended, per consuming project):

- **CM-CANONICAL-ROOT-CLARITY** — verifies (a) consumer's `CLAUDE.md` opens with the inheritance pointer (either `@import` or `## INHERITED FROM constitution/CLAUDE.md` heading), (b) the constitution submodule's three

files are present at the expected path, (c) no `## INHERITED FROM` block in the constitution submodule's own files (those ARE the source-of-truth, not consumers).

Propagation. Composes with §11.4.17 (universal-vs-project classification — §11.4.35 defines the file-layer split that §11.4.17 classifies INTO), §11.4.18 (script documentation discipline — same canonical-vs-consumer file-layer rule applies to docs/scripts/). Reading order: this anchor first, then §11.4.17 for the classification rule.

No escape hatch. Misfiling a rule between layers IS a §11.4.35 violation regardless of intent. The fix is `git mv + commit` per invariant 7, not “leave both copies in place to be safe” — duplicates silently diverge over time.

§11.4.36 — Mandatory `install_upstreams` on clone/add Mandate (User mandate, 2026-05-15)

Forensic anchor — verbatim user mandate (2026-05-15):

“Every Submodule or Git repository we add or clone MUST BE upstreams installed using Upstreamable utility which MUST BE available through exported paths of the host system (in `.bashrc` or `.zshrc`) using `install_upstreams` command executed from the root of the cloned (added) repository - only if in it is Upstreams or upstreams directory present with bash script files (recipes) for all repository's upstreams!”

Operative rule. Every time an operator or agent adds or clones a Git repository (`git clone`, `git submodule add`, `incorporate-submodule` per §11.4.31, etc.) under any consuming project of this Constitution, the post-clone procedure MUST be:

1. `cd` into the newly-cloned (or newly-added submodule's) working tree.
2. If the working tree contains an `upstreams/` directory (or legacy `Upstreams/` per the §11.4.29 transition) populated with one or more `*.sh` recipe files declaring upstream Git SSH URLs (the canonical declaration format defined by this constitution submodule's `Upstreams/*.sh` shape), the operator or agent MUST invoke `install_upstreams` from that directory.
3. `install_upstreams` is a host-system utility installed on the operator's PATH via `.bashrc` or `.zshrc` export. Implementation lives in the constitution submodule (`install_upstreams.sh`); operators alias / symlink it onto PATH once per host setup. The utility reads the recipe files, configures every declared upstream as a named git remote, and fans out origin push URLs across all declared upstreams.
4. If no `upstreams/` directory is present, no `install_upstreams` invocation is required.
5. If `upstreams/` is present but `install_upstreams` is not on PATH, the operator MUST install it (per the constitution submodule's setup docs) BEFORE making the clone usable.
6. Skipping step 2 when an `upstreams/` directory IS present is a §11.4.36 violation — the next push from that working tree will land on only one upstream, breaking §2.1 (Multi-upstream push is the norm).

Pre-commit attention. Before the first commit lands inside the newly-cloned working tree, the operator MUST verify that `install_upstreams` has been executed (when applicable). The quickest check: `git remote -v | grep -c push` reports the expected upstream count. If it reports 1 while `upstreams/*.sh` declares more, abort the commit and run `install_upstreams` first.

Automation. The constitution submodule's tooling (the future incorporate-submodule per §11.4.31, the existing `scripts/init-submodules.sh` patterns in consuming projects) MUST auto-invoke `install_upstreams` as part of any clone / add operation when the target tree has an `upstreams/` directory. Operator-explicit manual invocation remains supported.

Gate CM-INSTALL-UPSTREAMS-ON-CLONE. Pre-merge gate inspects every owned-by-us submodule pointer commit and verifies that: (a) if `upstreams/` (or legacy `Upstreams/`) is present in the submodule's tree, (b) the submodule's local checkout has at least as many configured push URLs as the declared recipe count. Paired mutation (§1.1): `remove one upstream from local origin --push config → gate FAILs`.

Composition. §11.4.36 composes with §2 (single-entripoint commit/push wrapper), §2.1 (multi-upstream push is the norm — this rule is its setup-time complement), §3 (submodule changes propagate through submodule commits first — requires multi-upstream parity), §9.2 / CONST-043 (no force-push — the multi-upstream wrap-up makes unforced parity easy), §11.4.17 (universal — every consuming project's clone procedure), §11.4.20 (subagent delegation when batch-cloning), §11.4.28 / CONST-051 (owned-submodule discipline), §11.4.29 / CONST-052 (`Upstreams/ → upstreams/` transition is exactly this rule's scope), §11.4.30 / CONST-053 (`.gitignore` ignores `upstreams/` only if the project explicitly decides not to track recipes, which is unusual — recipes ARE source of truth and SHOULD be tracked), §11.4.31 / CONST-054 (submodule-dependency- manifest works alongside `upstreams` recipes: `helix-deps.yaml` lists WHAT to add at parent root, `upstreams/` recipes list WHERE to push the resulting clones).

Classification: universal (per §11.4.17). No escape hatch. Severity-equivalent to §2.1 multi-upstream-push-violation at the clone-time setup layer — without the post-clone install, the working tree silently has a single-upstream blind spot.

§11.4.37 — Fetch-before-edit mandate (User mandate, 2026-05-15)

Forensic anchor — verbatim user mandate (2026-05-15):

“Make sure that `feedback_fetch_before_edit` memory rule is part of our constitution Submodule - the root Consitution, AGENTS.MD and CLAUDE.MD. Validate and verify that Proejct-Toolkit and all Submodules do inherit all of them! Follow the constitution Submodule documentation for details.”

Background. In multi-agent / multi-upstream codebases — where parallel Claude Code, Cursor, Aider, Codex, Gemini CLI, or human operator sessions may operate on overlapping scope — the local working tree’s state can lag behind the canonical upstream state by the time any given agent receives a task. Acting on stale local state produces three failure modes documented in the originating session (2026-05-15):

1. **Redundant work** — the agent re-does what a parallel session already finished, wasting tokens and operator review time. (In the originating incident, “Gitee removal” had already been committed upstream by a parallel agent 25 minutes earlier; the second agent’s `git remote remove gitee` was a no-op echo of work already done.)
2. **False confidence** — the agent reports completion of work that was already done by someone else, with no mechanical signal that their commit duplicates upstream history.
3. **Divergent history** — if the agent commits a parallel “completion” of already-done work, the result is two siblings of the same change, doubling the conflict surface for the next push attempt to multi-upstream remotes (§2.1).

The mandate. The FIRST git-touching action of any session, on any consuming project that participates in this constitution, MUST be:

```
git fetch --all --prune
git log --oneline HEAD..@{u}           # parent
git submodule foreach --recursive 'git fetch --all --prune --
    quiet'
```

If `HEAD..@{u}` is non-empty, the agent MUST integrate (ff-merge, rebase, or — if non-fast-forward — surface to operator per §11.4.4) BEFORE issuing any local edit, scanner run, or test cycle. The fetch step is non-negotiable even when the operator’s directive is phrased as “do X immediately” — the 30-second check prevents hour-long conflict reconciliation later.

Scope. Applies to:

1. The consuming project root.
2. Every owned submodule (per §11.4.28) — recursive.
3. The constitution submodule itself (§11.4.26 step 1 makes this explicit for constitution-side edits; §11.4.37 generalises it to ANY edit on the consuming project, not only constitution edits).
4. Any dependency cloned via `incorporate-submodule` (§11.4.31) or `git submodule add` (§11.4.36).

Anti-bluff invariant. The fetch+log check MUST produce captured evidence — the actual `HEAD..@{u}` output, even if empty. Skipping the check on the basis of “I just fetched” or “nothing could have changed in the last N minutes” is a §11.4.6 (no-guessing) violation: the remote state is not knowable without a fetch.

Composition. Composes with §2.1 (multi-upstream push — without the fetch, the agent can’t know which upstream has the canonical ref), §11.4.4 (test-interrupt-on-discovery — newly-fetched commits that contradict the planned work are exactly the “freshly discovered defect” that triggers cycle interruption), §11.4.6 (no-guessing — remote state requires fetch, not assumption), §11.4.20 (subagent delegation — parallel subagents MUST coordinate through fetched state, never assumed-local state), §11.4.26 (constitution update workflow — this rule generalises

§11.4.26 step 1 to ALL edits, not only constitution-file edits), §11.4.32 (post-constitution-pull validation — fetch-before-edit happens BEFORE the constitution-pull sweep §11.4.32 mandates).

Gates. Pre-build gate CM-FETCH-BEFORE-EDIT-AUDIT (when implemented in the consuming project) audits the most-recent commit range against the upstream HEAD at the commit's parent — if the parent ref was not the upstream HEAD at the time the commit was authored, FAIL. Paired mutation (§1.1): synthetic commit whose parent is N commits behind the then-current upstream HEAD — gate must FAIL.

Classification: universal (per §11.4.17). The rule has no project-specific assumptions — it applies to any multi-upstream / multi-agent codebase. No escape hatch. Severity-equivalent to a §11.4 PASS-bluff at the planning layer — acting on stale local state is the operational analog of asserting truth from unverified premises.

§11.4.38 — Installable-Asset Evidence Mandate (User mandate, 2026-05-17)

Forensic anchor — verbatim operator report (2026-05-17):

“app does not have a launcher icon anymore — the latest published app tester build(s). how come app passed anti-bluff checks?”

For any user-distributable build artifact (package, bundle, installer, or container image produced by the build pipeline and distributed to end users), tests and challenges **MUST** open the artifact and verify each user-visible asset is **present** and **non-degenerate**.

“User-visible asset” includes (non-exhaustive):

- Application icons for every density / resolution tier declared in the artifact metadata, including any platform-specific icon format (multi-resolution container, adaptive XML, symbol set, etc.) that the target OS uses when the minimum supported OS version requires it.
- Splash screens declared in the install metadata.
- Application name strings as declared in the install metadata.
- Any other asset whose absence causes the user to be unable to identify, launch, or interact with the installed application from the OS launcher / home screen / app drawer.

A PASS without opening the artifact and verifying the asset chain end-to-end is a §11.4 PASS-bluff, regardless of whether source files exist and tests otherwise pass. The specific failure mode this rule targets is: source file exists → build pipeline packages it → post-build checks pass at the source layer → artifact **ACTUALLY** produced with the asset stripped or misconfigured, and no gate ever opens the artifact to verify.

Required evidence: the anti-bluff challenge for a user-distributable artifact **MUST** produce per-asset PASS/FAIL lines showing: (a) the asset entry exists in the artifact package listing, (b) the asset is non-empty / non-degenerate (size ≥ platform-defined minimum OR format-validated), (c) where the OS's asset-resolution path involves indirection (alias, XML reference chain, density-qualifier override), the full chain is traced to the final rendered resource.

Consuming-project implementation: each consuming project ships one challenge script per artifact type that opens the produced artifact and verifies every declared user-visible asset. The challenge **MUST** run as part of the project's standard QA gate (equivalent of `make qa-all`).

Root cause of §11.4.38 introduction: a consuming project shipped multiple releases with the application icon absent on the target OS because: (1) the icon XML used a resource type that the OS resolves differently above a specific API level, (2) all pre-ship challenges verified source files only, none opened the packaged artifact. This is a pure §11.4 PASS-bluff — every test and challenge reported green while the end user saw no icon.

Classification: universal (per §11.4.17). No escape hatch. Severity- equivalent to a §11.4 PASS-bluff at the artifact-packaging layer. Composes with §11.4.1–§11.4.5 (evidence requirements), §11.4.25 (full automation coverage), §11.4.27 (no fakes beyond unit tests — source-layer checks of a distributable artifact are the distributable- layer analog of unit-test-only coverage). See Constitution §11.4.38 for the full mandate.

§11.4.39 — Per-Feature On-Device End-User Validation Mandate (iter-76, 2026-05-17)

Every user-facing feature in a consuming project **MUST** have at least one HelixQA scenario that exercises the feature from cold-launch with **positive runtime evidence**:

- A screenshot or screen recording captured at a named checkpoint.
- At least one assertion of type `screenshot`, `accessibility_count`, `accessibility_node`, `ocr_text`, `pixel_histogram`, or `editor_state`.
- Evidence files written to a discoverable output directory so every PASS is cross-verifiable after the fact.

Scenarios **MUST** be RE-EXECUTED on every release candidate to catch cross-iteration regression. A scenario authored but never re-run on subsequent iterations degrades to a metadata-only gate — a §11.4 PASS-bluff at the regression layer.

Authoring rules (consuming project sets project-specific paths):

1. Scenarios are stored in a canonical bank directory designated by the consuming project (e.g. `<banks-root>/feature-coverage/`).
2. Each scenario YAML **MUST** include: `name`, `version`, `metadata`, `platforms`, `test_cases`, and at least one step with `evidence_required: true` and a specific `evidence_type`.
3. A **coverage matrix** document **MUST** exist listing feature × iteration × scenario. Matrix rules: (a) new iteration **MUST** add ≥ 1 scenario per new user-facing feature; (b) prior scenarios **MUST NOT** be deleted (mark as `status: retired` if feature is removed); (c) all non-retired scenarios **MUST PASS** for ship-ready.
4. A **static gate** (challenge script) **MUST** assert the scenario count $\geq N$ (where N is the number of user-facing features registered in the coverage matrix) and that each YAML contains a `evidence_type: assertion`.
5. **Portable host tooling:** evidence-validation scripts **MUST** use POSIX-portable file-size helpers (e.g. `wc -c < file`) instead of OS-specific commands (e.g. GNU `stat -c%5`) to avoid silent 0-byte reports on BSD/macOS hosts. A regression challenge **MUST** verify the portable helper on the actual host.

iOS / platform-deferred pattern: scenarios written platform-agnostically with the target platform listed in a comment deferral (e.g. # iOS: deferred – tracker #X) are valid; add the platform to the `platforms:` list when the automation toolchain becomes available — no structural change to the scenario is needed.

Classification: universal (§11.4.17). Composes with §7.1, §11.4 (anti-bluff), §11.4.25 (full automation coverage), §11.4.27 (no fakes), §11.4.38 (artifact evidence). No escape hatch — a feature without an on-device scenario is NOT covered per §11.4.25 invariant 2. See Constitution §11.4.39 for the full mandate.

§11.4.40 — Full-suite retest before release tag mandate (User mandate, 2026-05-17)

Forensic anchor — verbatim user mandate (2026-05-17):

“Please, do complete retests with all existing tests (multi-hours efforts) when all new workable items are done, fixed, polished and verified. Time is essential! We should already have this in our (root) Constitution, CLAUDE.MD and AGENTS.MD.”

Operative rule. A release tag (any `vX.Y.Z / X.Y.Z-suffix` identifier MUST NOT be created until a **COMPLETE retest with ALL existing tests** has been executed on a clean baseline AFTER every workable item in the batch is done, fixed, polished, and individually verified. A “spot-check” retest that runs only the tests directly touched by the batch is FORBIDDEN — it misses interaction defects between the batch’s fixes and previously- stable code paths.

The complete retest comprises:

1. **Pre-build verification full sweep** — every gate in the project’s pre-build script runs; 0 FAIL required.
2. **Post-build verification full sweep** — every gate in the project’s post-build script runs against the freshly-assembled image; 0 FAIL required.
3. **On-device 4-phase cycle** (e.g. `test_all_fixes.sh` or project equivalent) on **EVERY owned device** (the full topology at that point in time). Each phase (Immediate Fresh Flash / After Reboot / After Factory Reset / After Final Reboot) MUST complete; no phase may be skipped. Phase summaries captured.
4. **Meta-test full mutation sweep** — every paired mutation in `meta_test_false_positive_proof.sh` (or project equivalent) executed; every gate proven non-bluff.
5. **Test bank full sweep** — if a Challenge-driven test bank exists, every Challenge in the bank covered by a full QA session (not a sub-set).
6. **Issues.md / Fixed.md state audit** — no item silently demoted; every Reopened item has §11.4.34 Reopened-Details; every closure has captured-evidence per §11.4.5; no `Status:` value outside the §11.4.15 closed-set.
7. **CONTINUATION.md sync check** — per §12.10, document reflects current state at the moment of tagging.

Time is essential. A complete retest is typically a **12–48 hour elapsed effort** depending on parallelism, device count, and meta-test mutation count. This is NOT optional and NOT abbreviated. Operators should plan release cadence with this duration in mind, NOT skip the retest to ship faster. Skipping the retest is the exact “tests passed but feature broken” failure mode §11.4 specifically prohibits.

Composition with §11.4.4. Per-fix retest (`test_all_fixes.sh` run after each individual fix lands) is STILL mandatory per §11.4.4. §11.4.40 is the **additional final integrity check** that runs after the BATCH is complete — catching interaction defects that individual per-fix retests cannot see in isolation.

Composition with §11.4.7. The complete retest is the authoritative captured-evidence baseline for any closure of items present in the batch. If a §11.4.7-promoted item PASSED its per-fix retest but FAILs the full-suite retest, the closure MUST be reverted and the item moved back to In progress — captured-evidence-contradicts under same-conditions per §11.4.7.

Composition with §11.4.39. Per-feature on-device end-user validation runs as part of step 3 (on-device cycle) — every feature’s wrapper-test fires during the full-suite retest. The two mandates are complementary: §11.4.39 covers individual feature validation breadth; §11.4.40 covers release-time integration depth.

Gate CM-FULL-SUITE-RETEST-MANDATE. Pre-tag gate inspects the release-candidate state and verifies that within the last 72 hours of the candidate commit, evidence exists of (a) full pre-build + post-build run, (b) on-device 4-phase cycle on every owned device, (c) meta-test full sweep, (d) Issues.md/Fixed.md audit pass. Evidence captured in `docs/changelogs/<tag>.md` per §11.4.4(c) requirements. Paired mutation (§1.1): strip Full-suite retest evidence block from a changelog → gate FAILs.

Classification: universal (per §11.4.17) — every consuming project’s release procedure. No escape hatch — there is no `--skip-full-retest` flag, no `--quick-release` mode. Operators who feel time-pressured to skip the full retest should instead delay the release until the retest completes; shipping unverified code is worse than delayed shipping.

§11.4.41 — Pre-Force-Push Merge-First Mandate (User mandate, 2026-05-17)

Forensic anchor — verbatim user mandate (2026-05-17):

“make sure we bring everything from branches to our side before forc push is done! Afer everything is safely and fully merged and all potential conflicts (if any) resolved, then do force push! make sure nothing isnlost, broken or corrupted on bith sides! add these rules in our root Constitution, CLAUDE.MD, AGENTS.MD (constitution Submodule) if itnis not added already! Extremely important rules and mandatory constraints we MUST HAVE and fully respect!”

Operative rule. Any force-push (`git push --force`, `git push --force-with-lease`, `git push +<ref>`, equivalent history-rewriting operation on any remote) authorised under §9.2 / CONST-043 MUST be preceded by a mechanical 4-step merge- first pipeline that brings every remote-side commit into the local tree, resolves every conflict carefully, and verifies nothing is lost or corrupted on EITHER side BEFORE the overwriting push is executed.

The 4-step pipeline (mandatory, in order):

1. **Fetch every remote-side reference.** Run `git fetch --all --prune --tags` against every configured remote (origin + every upstream). Capture the output.

2. **Integrate every divergent commit locally.** Compute `git log --oneline HEAD..<remote>/<branch>` for every remote. For every non-empty range, integrate the remote commits into the local tree via the appropriate strategy:
 - **git rebase <remote>/<branch>** when the local divergent work is a strict superset that should land on top, OR
 - **git merge <remote>/<branch>** when the local + remote histories are independent additions that both deserve preservation, OR
 - **operator-confirmed cherry-pick** when remote contains a subset of commits already present locally in different form.
3. **Audit the integrated tree.** Verify (a) no conflict markers anywhere in the working tree (`grep -rn '^<<<<<< \|^===== $\|^>>>>>> '` . returns empty across all governance + source + test files — third-party docs containing the markers as illustrative content excepted), (b) no file silently dropped (`git diff --stat HEAD@{1} HEAD` shows only expected additions), (c) every previously-passing test still passes on the integrated tree (per §11.4.4 + §11.4.40 baseline), (d) every captured- evidence artifact from the pre-merge state still validates.
4. **Execute force-push.** Only after steps 1-3 produce captured evidence of clean integration: `git push --force-with-lease <remote> <ref>` (NEVER `--force` without `--with-lease` unless explicitly authorised per §9.2 sub-clause 6 for a specific remote where lease semantics are unavailable). One force-push event per CONST-043 authorisation — no batch authorisation across multiple force-pushes.

Three failure modes prevented:

1. **Remote-side content loss.** Force-push without prior fetch silently overwrites commits a parallel session, sibling agent, or co-developer landed on the remote. The lost work is recoverable from the remote's reflog only within its TTL window — typically 30-90 days — but the audit trail (PR comments, CI evidence, governance signatures) is lost permanently.
2. **Stale-state act on remote.** Force-push from a divergent local state that was branched from a remote tip N hours ago treats the remote-side N hours of work as nonexistent. The operator's `--force-with-lease` safety check catches the literal SHA mismatch but ONLY when invoked AFTER a fetch that updated the lease reference; without step 1, the lease check is reading stale local refs.
3. **Conflict-driven corruption.** A merge that completes with unresolved markers in the tree (one half kept, other dropped silently because the marker was inside a comment block, OR markers committed verbatim as in CONST-049-prerequisite incidents observed 2026-05-17 on helix_qa + containers) lands broken governance / source on the force-push target. Step 3's explicit conflict-marker `grep` is the mechanical guard.

Two-gate composition with CONST-043. §11.4.41 does NOT relax CONST-043's operator-approval requirement — it adds a SECOND gate on top:

- **Gate A (CONST-043).** Operator types explicit per-operation force-push authorisation in the conversation.
- **Gate B (§11.4.41).** Agent executes the 4-step merge-first pipeline, captures evidence of clean integration, presents evidence to operator BEFORE the force-push.

Both gates required. CONST-043 alone authorises a force-push that loses remote work; §11.4.41 alone risks force-pushing without operator awareness. The two together produce a force-push that is BOTH operator-authorised AND remote-safe.

Verification artefact. Every §11.4.41-governed force-push emits a docs/change logs/<tag>.md “Force-push merge-first audit” section containing: (i) git fetch output, (ii) per- remote HEAD..<remote>/<branch> log before integration, (iii) integration strategy chosen per remote with rationale, (iv) post-integration conflict-marker scan output (must be empty), (v) post-integration test suite delta (must show only expected changes), (vi) the --force-with-lease push output with lease SHA evidence, (vii) CONST-043 authorisation quote from the conversation.

Cascade requirement (per CONST-047). This anchor (verbatim or by §11.4.41 / CONST-061 ID reference) MUST appear in every owned submodule’s CONSTITUTION.md, CLAUDE.md, and AGENTS.md. Severity-equivalent to a §11.4 PASS-bluff at the remote-data-integrity layer. Gate CM-FORCE-PUSH-MERGE-FIRST walks docs/change logs/<tag>.md “Force-push” entries for the 7 audit elements; paired mutation strips any element and asserts gate FAILs.

Classification: universal (per §11.4.17) — every owned project executes force-pushes against its remotes; the merge- first discipline is reusable verbatim. No escape hatch — the operator-pressure escape (“just force-push, we’ll fix it later”) is the exact failure mode this anchor closes.

Composes with: §9.2 (data-safety hardlinked backup), §11.4.4 (test-interrupt-on-discovery — broken integration triggers rollback), §11.4.6 (no-guessing — every step’s outcome captured, not assumed), §11.4.26 (constitution-submodule update pipeline — that mandate is the per-submodule specialisation of §11.4.41 for governance-files-only commits), §11.4.32 (post-pull validation — the audit step’s mechanical companion), §11.4.37 (fetch-before- edit — step 1 enforces it for the force-push specifically), §11.4.40 (full-suite retest — step 3’s test-evidence prerequisite), CONST-043 (per-operation operator approval), CONST-047 (recursive cascade). ### §11.4.42 — Iteration-discipline mandate (User mandate, 2026-05-18)

Forensic anchor — verbatim user mandate (2026-05-18):

“Are we clear about working iterations and sorting priorities for testing and fixing? We first fix top and middle critical priorities → We flash and test all of these → If nothing is broken and no new issues are reported by users or found we run full system testing (many hours for everything) → if anything is still broken or new issues are reported we get back to fixing → Cycle repeats. Make sure this is absolutely clear and mandatory foundation of our root (constitution Submodule) Constitution, AGENTS.MD and CLAUDE.MD! We MUST WORK in this manner until otherwise has been told.”

Operative rule. Project work proceeds in priority-ordered iteration cycles. Each cycle is a closed loop of five steps that MUST run in order; advancing past any step before its acceptance condition is met is a §11.4 bluff equivalent to skipping the step entirely.

The five-step cycle:

1. **Priority-ordered fix selection.** Open the project's issue tracker, filter to items whose `**Status:**` is one of `Queued` | `In progress` | `Reopened` | `Operator-blocked` (per §11.4.15 closed-set), sort by severity DESC then intra-group criticality DESC (per §11.4.12 sort order), select ONLY items classified `Top critical` or `Middle critical` for the current batch. Low items are explicitly DEFERRED until the critical batch is closed.
2. **Batch implementation.** Land source-side fixes for the selected batch per §11.4.9 (batch-source-fixes-before-rebuild) — every non-runtime-dependent fix lands BEFORE the next rebuild, every fix gets the §11.4.4 four-layer coverage (pre-build gate + post-build gate + on-device test + paired mutation).
3. **Smoke-test gate.** After flash, run the SMOKE bundle:
 1. anti-bluff baseline (`meta_test_false_positive_proof.sh` for gates touched by this batch), (b) every `test_*.sh` whose name matches a fix in this batch, (c) the critical-path regression probe (boot + launcher + WiFi + audio + secondary-display routing core), (d) the §11.4.39 per-feature on-device validation for any feature touched. Estimated runtime: <30 minutes on the parallel device fleet. Smoke MUST emit zero FAIL and zero unclassified WARN.
4. **Full-system-test gate.** ONLY if (a) smoke is GREEN AND (b) no new operator/user report has surfaced during the batch's validation window, run the §11.4.40 complete retest (all 7 steps including 12–48 h on-device 4-phase cycle on every owned device). If smoke FAILED OR a new report arrived, the full-system test is FORBIDDEN; loop back to step 1 with the new evidence added to the priority queue.
5. **Release-readiness or loop.** If full-system test is GREEN AND no new report arrived during it, the batch is release-ready (the §11.4.40 tagging procedure may proceed pending operator authorization). Otherwise loop back to step 1.

Priority taxonomy (binds existing severity convention). `Top critical` = severity C (Critical per Issues.md convention) AND intra-group criticality 5. `Middle critical` = severity C with intra-group 1–4 OR severity M (Major) with any intra-group score. Low = severity L per the existing [C/M/L] taxonomy documented in §11.4.12. Severity WARN items are treated as Low for iteration-discipline purposes.

Cycle repeats. The five-step loop continues, fed by the priority queue, until the operator explicitly authorizes a release. There is no upper bound on iteration count; “we’ve cycled enough” is not a Constitution-recognised stopping condition. Time is essential per §11.4.40 BUT the cycle MUST NOT be truncated to ship faster — that is exactly the bluff §11.4 forbids.

Composition. §11.4.4 (per-fix retest inside step 2) + §11.4.7 (closures require same-conditions positive evidence; step 4 is authoritative baseline) + §11.4.9 (source-side batching inside step 2) + §11.4.34 (Reopened items attribute By: AI / User) + §11.4.40 (the multi-hour full-suite retest IS step 4 — §11.4.42 is the meta-loop conductor binding the instruments).

Anti-bluff coupling. Per §11.4.2 + §11.4.5, every smoke-test and full-system-test PASS MUST carry positive captured evidence of user-visible behaviour. Tests AND HelixQA Challenges bound equally — a Challenge that scores PASS without applicable analysis is a §11.4 PASS-bluff.

No escape hatch. No `--skip-priority-batch`, `--skip-smoke`, `--full-suite-only`, or `--release-without-loop` flag exists. Subagents performing autonomous work default to the §11.4.42 path. Operators who feel time-pressured to short-circuit the loop should add capacity (more devices in parallel, more agent slots) rather than skip steps. Shipping under-validated code is worse than delayed shipping.

Gate CM-COVENANT-114-42-PROPAGATION mirrors CM-COVENANT-114-40-PROPAGATION 1:1 — every CLAUDE.md / AGENTS.md in the covenant file set carries the §11.4.42 anchor. Paired mutation strips the anchor literal from one consumer file → gate FAILs.

Classification: universal (per §11.4.17) — every consuming project’s iteration workflow.

§11.4.43 — TDD-Fix-Discipline Mandate (User mandate, 2026-05-18)

Forensic anchor — verbatim user mandate (2026-05-18):

“Make sure you do validate and verify every fix! Make sure you first do the fix on live device using ADB if that is possible, once fix is confirmed, make sure that it is fully applied into our codebase ... It is important to note that we shall start by creating test which confirms the issue (fails as expected), and ones the fix is produced it becomes positive - it does pass with success! No bluff policy is mandatory and nothing can be bluffed! Every step of this process MUST BE 100% bluff-free and bluff-proofed!”

Operative rule. Every fix MUST follow the 5-step TDD-fix workflow:

1. **RED — Failing test FIRST.** Before any code change, author a test (or extend an existing one) that exercises the user-visible path of the reported defect and FAILs for that defect specifically. Per §11.4.1 the FAIL MUST be a real product defect (not a script-bug induced FAIL). Per §11.4.2 the FAIL run MUST produce captured evidence (`logcat` / `dumpsys` / `screenrecord` / `dmesg` / `/sys` snapshot / sink-side probe). The test identifier MUST cite the §-letter or Fix-# under repair. A RED step that “happens to pass because the underlying probe lifecycle isn’t wired” is itself a §11.4 PASS-bluff.
2. **LIVE-ADB-PROBE — try the fix on the running device first (when feasible).** For mutable surfaces — `setprop persist.*`, `settings put`, `pm clear`, `am force-stop`, `am start`, `cmd wifi`, `cmd media_session`, `ad-hoc / sys writes`, `boot- script edits pushed to /vendor/bin/`, `test-fixture replacement` — the operator MUST attempt the fix on the running device via `adb shell` first. The fast `adb-loop` (~5 min per attempt) replaces the `build-flash-test loop` (~6 h per attempt) for confirmation of the hypothesis. The live-probe IS NOT the fix; it is the EXPERIMENT that proves the hypothesis is correct before the source-side investment.

Live-probe is INFEASIBLE (exception, must rebuild) for: kernel changes, AOSP framework (`frameworks/base/`), hardware HAL changes, `vendor/` partition contents, `init.rc` service definitions, `sepolicy`, `ro.*` build properties (immutable post-boot), AOSP-build-time XML / resource overlays, `Android.bp/`

Android.mk changes, signed-system-app contents (LOCAL_CERTIFICATE := platform). Each exception MUST be cited in the commit message as LIVE_PROBE_INFEASIBLE: <reason>.

3. **GREEN — apply the fix to source code.** The source patch MUST achieve the same effect on the device as the live probe. Per §11.4.9 the source change is batched with other source-side fixes where appropriate. Build via the project's containerized build (§12.9). Flash via the project's flash script. The bytes on the assembled image MUST land where the post-build gate expects (per §11.4.4(b) four-layer requirement).
4. **VERIFY — re-run the RED test, must now PASS.** Per §11.4.7 demotion-evidence: the PASS MUST come from the SAME conditions that initially FAILED (same device, same firmware NOW carrying the fix, same load profile, same cycle position). Per §11.4.5 the PASS MUST carry positive captured evidence (presence + correctness — RMS amplitude, fprobe channel count, frame count, OCR text, sink codec state, etc.). Per §11.4.2 video tests MUST cross-check via the recording-analyzer. Per §11.4.42 a reliability check at the iteration count (typically 10 iterations) MUST PASS all iterations — a single intermittent FAIL keeps the item in In progress per §11.4.7.
5. **DOCUMENT — update every relevant doc in the SAME commit.** Per §11.4.4(c): docs/Issues.md → docs/Fixed.md migration with type-aware closure vocabulary per §11.4.33 (Bug → Fixed, Feature → Implemented, Task → Completed); the project's CLAUDE.md Applied Fixes Reference row; per-version docs/changelogs/<tag>.md entry; affected user-facing guides under docs/guides/; HelixQA bank entry per §11.4.4(b); docs/CONTINUATION.md per §12.10; project memory file if non-obvious.

Composition. §11.4.43 is the workflow that ENACTS the existing covenant. Composes with: §11.4.1 (the RED FAIL is a real defect), §11.4.2 (captured evidence at both RED and VERIFY), §11.4.3 (topology dispatch in the RED test), §11.4.4 (four-layer coverage at GREEN), §11.4.5 (quality analysis at VERIFY), §11.4.6 (no guessing during the LIVE-PROBE hypothesis statement — UNCONFIRMED: until probe proves it), §11.4.7 (demotion-evidence at VERIFY — same conditions), §11.4.8 (research citation precedes step 1), §11.4.9 (GREEN is batched), §11.4.40 (full-suite retest is the release-tag-time additional check), §11.4.42 (10-iteration reliability loop IS step 4's inner cycle).

Gate CM-COVENANT-114-43-PROPAGATION. Pre-build gate verifies the §11.4.43 anchor is present in every CLAUDE.md / AGENTS.md across the covenant file set. Paired mutation strips the anchor literal from one consumer file → gate FAILs.

Classification: universal (per §11.4.17) — every consuming project's fix workflow.
No escape hatch. No --skip-red-test, --no-live-probe, --skip-verify, --skip-document flag exists. "I'll add the test after the fix" is the exact PASS-bluff pattern §11.4 forbids because the test authored after the fix demonstrates only that the test agrees with the fix, not that the test catches the bug.

§11.4.44 — Document Revision Header Mandate (User mandate, 2026-05-18)

Forensic anchor — verbatim user mandate (2026-05-18):

“Add this change / improvement as background work in parallel with our mainstream work. All documents (Issues, Issues_Summary, Fixed, Continuation doc and all others MUST HAVE revision number! Revision number with information such as: last date and time when document was modified MUST GO at the top of the document below the main H1 title. Add all relevant information about the document as well. However, last date and time document has been modified and revision number are MANDATORY! Revision number shall be the whole number 1, 2, 3, etc. Document all this in our root (constitution Submodule) Constitution, AGENTS.MD and CLAUDE.MD as mandatory rule(s) and constraints we MUST HAVE! No bluff is allowed of any kind!”

The mandate. Every IN-scope tracked Markdown document MUST carry a header block directly below the H1 title containing two MANDATORY fields:

- ****Revision:**** N where N is a monotonic positive integer (1, 2, 3, ...). Never decremented. Never reset on rewrite. A freshly- created IN-scope document starts at Revision 1.
- ****Last modified:**** YYYY-MM-DDTHH:MM:SSZ (ISO 8601 UTC).

Optional encouraged fields: ****Description:****, ****Authority:****, ****Maintainer:****, ****Scope:****, ****Auto-generated-from:**** (auto-generated docs only), ****Status:**** (plan docs only).

Header format (canonical):

Document Title

```
**Revision:** 47
**Last modified:** 2026-05-18T13:42:00Z
**Description:** Open / in-flight bug + work tracker
**Authority:** Constitution §11.4 covenant
**Maintainer:** Operator + AI loop per §11.4.42
```

<content...>

YAML front-matter and HTML-comment headers are FORBIDDEN — they are invisible to humans scrolling the raw Markdown and break the “all relevant information about the document” visibility clause.

IN scope (revision header MANDATORY): docs/Issues.md, docs/Issues_Summary.md, docs/Fixed.md, docs/Fixed_Summary.md, docs/CONTINUATION.md, docs/guides/**/*.*.md, docs/research/**/*.*.md, docs/scripts/**/*.*.md, docs/changelogs/**/*.*.md, docs/superpowers/plans/**/*.*.md, docs/hardware/**/*.*.md, and every other tracked Markdown under docs/.

OUT scope (revision header NOT required): CLAUDE.md / AGENTS.md (every layer — already version-tracked via VERSION file + inheritance pointer per §11.4.35; adding a per-file revision would duplicate authority and obscure the canonical-root contract); README.md / CONTRIBUTING.md / LICENSE / NOTICE / VERSION / OWNERS (standard metadata); rendered HTML/PDF artifacts (revision inherited from source Markdown); auto- generated docs that already carry an embedded generation-time stamp.

Auto-bump tooling: - `scripts/doc_revision_bump.sh <file>` — manual entry point, idempotent (no-op if Last-modified within 5 s of date `-u`). - `.git/hooks/pre-commit` — automatic entry point; walks staged IN-scope docs, calls the bump for each, re-stages. - `scripts/testing/sync_issues_docs.sh` — composes the bump for `Issues_Summary` + `Fixed_Summary` after regeneration. - `scripts/commit_docs.sh` — calls the bump before staging so operators who skipped the hook installation are still covered.

Composition with §12.10. `CONTINUATION.md` already carries `Last updated:` per §12.10. §11.4.44 adds `**Revision:** N` directly under `H1` + reuses the existing `Last updated:` line as the §11.4.44 `Last modified:` line. One source of truth, two pointers — no duplication.

Composition with §11.4.12. `Issues_Summary.md` / `Fixed_Summary.md` inherit the revision number of their source-of-truth Markdown at generation time; they are NOT incremented independently.

Composition with §11.4.23. HTML colorizer preserves the revision header in the colored HTML — pandoc renders the bold lines as styled `<p><strong>...</strong></p>` blocks which the colorizer leaves alone (only `Issues-table` rows are mutated).

Gates: `CM-DOC-REVISION-HEADER-PRESENT` walks IN-scope docs and asserts both mandatory fields appear within 15 lines below the `H1` title; `CM-COVENANT-114-44-PROPAGATION` asserts the anchor literal is present in every covenant file. Paired mutations strip the `**Revision:**` line OR the anchor literal → gates FAIL.

Classification: universal (per §11.4.17). **No escape hatch** — no `--skip-revision-bump`, `--no-header`, `--allow-missing-revision` flag exists. Operators who feel pressured to skip the bump should land smaller commits, not bypass the discipline.

§11.4.45 — Integration-Status-Doc Maintenance Mandate (User mandate, 2026-05-18)

Forensic anchor — verbatim user mandate (2026-05-18):

“Make sure this Markdown document is regularly updated to reflect current up to date state of full integration ... Whole integrations MUST be regularly updated and retested! Keep the Markdown document ALWAYS up to date and ALWAYS exported into PDF and HTML! Make sure document is ALWAYS in sync! Add all this into our Constitution, `AGENTS.MD` and `CLAUDE.MD`.”

Generalisation note. §11.4.45 is the generic form of §12.10 (`CONTINUATION.md`) applied to every domain integration. §12.10 binds the canonical handoff document specifically; §11.4.45 generalises the sync-and-evidence pattern to every non-trivial integration (Dolby/Arvus, NanoKVM, Firebase, Widevine L1/L3, Play Protect, Netflix, Cuttlefish, Chromecast, WiFi chip, ES8388 codec, Sonos, AC-3, Google Stack, HiFi audio, TV apps, etc.). Without this generalisation each new integration re-invents the same sync wrapper, revision-header discipline, captured-evidence requirement, and operator-blocked surface — duplicating effort and risking drift.

Operative rule (10 sub-points — every Status.md MUST hold ALL).

Every integration-status document at docs/<domain>/<integration>/
Status.md:

1. MUST exist when a domain integration is non-trivial (more than one fix / test / gate landing for that integration).
2. MUST carry the §11.4.44 revision header directly below the H1 title (Revision + Last modified + Description + Authority + Maintainer + Scope).
3. MUST be auto-synced (HTML + PDF) on every related-test-cycle completion AND every fix touching the integration. “Auto-synced” means the wrapper from sub-point 5 is invoked; manual regeneration is forbidden as the sole path because operators will forget.
4. MUST be auto-colored per §11.4.23 — Status column cells colored by lifecycle state, Type cells colored by Type, operator-blocked rows visually distinct.
5. MUST have a sync wrapper invocable as `bash scripts/testing/sync_integration_status.sh <Status.md>` OR a per-integration thin shell that delegates to the generic wrapper. The thin shell exists so operators can type a short command per integration without remembering paths.
6. MUST live under docs/<domain>/<integration>/ directory structure — consistent layout, discovery by glob remains O(1).
7. MUST include a captured-evidence-driven status table per §11.4.5 — every status claim cites the test log path, recording file path, or sink-probe report that backs it. Status claims without evidence are §11.4 PASS-bluffs.
8. MUST distinguish status values per §11.4.6 / §11.4.7: PASS / FAIL / SKIP / PENDING_FORENSICS / OPERATOR-BLOCKED — closed vocabulary, mechanically distinguishable, no implicit values.
9. MUST list operator-blocked items at the top of the document (just below the revision header) so an operator scanning the Status.md finds action items in O(1) — not buried 200 lines deep in a chronological evidence log.
10. MUST be referenced from docs/CONTINUATION.md §3 (Active work) when any item in the Status.md is non-terminal (Queued / In progress / Reopened / Operator-blocked) — composes with §12.10 so the canonical handoff document points at the integration-status doc rather than re-stating its contents.

Gates:

- CM-COVENANT-114-45-PROPAGATION — anchor text propagated to every CLAUDE.md / AGENTS.md across the covenant file set.
- CM-AF-INTEGRATION-STATUS-DOCS — discovers all docs/**/Status.md via glob, verifies each carries the §11.4.44 header, has a sync wrapper, HTML+PDF exports are mtime-current vs the source, colorization applied (cell-status class present in the colorized HTML).

Paired meta-test mutations (§1.1 compliance): propagation strip → propagation gate FAILs; delete `**Revision:**` from a Status.md → status-docs gate FAILs; touch Status.md without re-syncing → mtime mismatch FAILs; inject likely outside UNCONFIRMED: block → existing CM-NO-GUESSING-MANDATE catches it (composition catch).

Composition. §11.4.5 (captured evidence in every row), §11.4.6 / §11.4.7 (status vocabulary closed-set), §11.4.12 (sync wrapper pattern reused for HTML+PDF re-export), §11.4.13 (sink-side captured evidence is a specific instance of the generic

rule), §11.4.15 (status values), §11.4.22 (commit_docs.sh wrapper), §11.4.23 (colorizer extends to Status.md HTML), §11.4.44 (every Status.md carries the revision header), §12.10 (CONTINUATION.md references Status.md paths in §3).

No escape hatch — no `--skip-status-sync`, `--no-revision-bump-on-status`, `--allow-stale-html` flag. Operators who feel pressured to skip the sync should land smaller, more frequent integration commits rather than batching status updates into one giant push.

Classification: universal (per §11.4.17).

§11.4.46 — Validate-recent-work-before-post-flash-tests mandate (User mandate, 2026-05-18)

Forensic anchor — verbatim user mandate (2026-05-18):

“We see that on newly flashed devices we execute all post flash tests. Make sure we do execute them all after we validate and verify all recent work from the Issues and Continuation docs first! Once all new / latest work is fully validated and verified, no new issues found, nothing is discovered broken or faulty, then we run all post-flash tests! If validation and verification of recent work fails for any reason, and we must go back to fix anything that popped up we do not have to run post-flash tests! We MUST save time! All post-flash tests can be executed only after recent work on features and issues is fully validated and confirmed with complete anti-bluff policy enforced! Add this so it is absolutely clear into our root (constitution Submodule) Constitution, AGENTS.MD and CLAUDE.MD! Start following all these rules IMMEDIATELY! Iff required abort all post-flash testing if it is in progress!”

Operative rule. After every device flash, the orchestrator MUST first run a recent-work validation pass before any full post-flash suite (`test_all_fixes.sh` or project equivalent). The validation pass enumerates recent work from three docs (Issues.md / Fixed.md / CONTINUATION.md §3), maps each recent item to its on-device test binding, runs ONLY those tests, classifies each result per §11.4.6, and demands 100% green before the full suite is permitted.

Definition of “recent work” (closed-set):

- docs/Issues.md actionable headings with `**Status:**` ∈ {In progress, Ready for testing, Reopened} (per §11.4.15 closed-set).
- docs/Fixed.md headings closed within `--max-age-days N` (default 7).
- docs/CONTINUATION.md §3 “Active work” sub-headings listed as IN PROGRESS or BLOCKED.

Recent-work test binding. Each recent item’s on-device test is discovered via §-letter or Fix # match against the project’s `test_*.sh` filenames OR an explicit `Test:` line in the item heading. Items without an on-device test are §11.4.4 four-layer-coverage violations (a fix without an on-device test is forbidden) — flagged as VALIDATION-PASS-BLOCKED until the test lands.

Authorization sequence (machine-enforced):

1. `scripts/testing/recent_work_validate.sh --device <serial>` runs; writes `/data/local/tmp/.recent_work_validated` on the device IFF all targeted tests return PASS with captured evidence.
2. The full suite (`test_all_fixes.sh`) checks for the marker file at entry; absent or stale (older than the last flash boot epoch) → refuses to proceed with exit code 11.
3. Marker is invalidated automatically on reboot — the orchestrator stores the device boot epoch in the marker and re-validates.

Anti-bluff during validation. Each recent-work item that landed a fix in this batch MUST have a paired §11.4.43 RED test (captured before the fix) AND the same test now GREEN (captured after the fix). A GREEN with no prior RED is itself the bluff §11.4 forbids — the gate FAILs.

Honest classification of validation FAILs. Per §11.4.6: PRODUCT defect (real regression — block full-suite, surface to operator), TEST-INFRA defect (test-script bug — fix per §11.4.1 at source layer, re-run), BLUFFY-THRESHOLD (e.g. SLO floor wrong — re-baseline), or UNCONFIRMED (not reproducible in isolation — tag per §11.4.7 PENDING_CYCLE_RETEST, NEVER silently demote).

Composition. §11.4.4 (STOP-on-discovery) + §11.4.6 (no- guessing) + §11.4.7 (demotion-evidence) + §11.4.40 (full-suite gate) + §11.4.42 (iteration discipline — smoke-test ≡ recent-work- validation) + §11.4.43 (RED test for each recent-work item IS the validation test) + §11.4.44 (revision header determines recency) + §12.10 (CONTINUATION.md source-of-truth for §3 Active work).

Gates:

- CM-COVENANT-114-46-PROPAGATION — anchor literal present in every CLAUDE.md / AGENTS.md across the covenant file set.
- CM-AF-RECENT-WORK-VALIDATION-GATE — asserts `scripts/testing/recent_work_validate.sh` exists + executable + sources the anti-bluff library + the full-suite orchestrator contains the entry-gate check.
- CM-AF-VALIDATION-ARTIFACT-FILE — asserts the marker file path `/data/local/tmp/.recent_work_validated` is literal-identical across helper + orchestrator + propagation-block + this Constitution section.

Paired mutations strip the anchor / remove the entry-gate / flip the marker path → respective gate FAILs.

No escape hatch. No `--skip-validation`, `--full-suite-always`, `--ignore-recent-work` flag is permitted. The full-suite-without- validation pattern is the precise “tests passed but feature broken” failure mode §11.4 specifically prohibits. The 60–90 min validation pass is ALWAYS faster than 4–6 h of full-suite chasing a defect.

Classification: universal (per §11.4.17).

§11.4.47 — Firebase Data Review Mandate (User mandate, 2026-05-18)

Forensic anchor — verbatim user mandate (2026-05-18):

“We MUST regularly before every bigger working round check Firebase information gathered so far: Crashlytics (all fatals, non-fatals and ANRs), Analytics data and Performance data. For anything problematic depending on severity proper workable items (issues) MUST be created (Issues, Issues_Summary docs) with full references (links) to original data on Firebase. We MUST make sure we do not create for same issues multiple entries, so we MUST distinguish between same issue manifested in different forms (stacktraces)! Everything MUST BE comprehensive and in-depth level of details so any changes we do (or fixes) do solve the original root problems!”

The mandate. Before every “bigger working round” (pre-build, pre-flash, pre-tag, daily, post-deployment burn-in) the operator/ loop MUST execute the Firebase review pass via `scripts/firebase/review_round.sh`. The pass queries Crashlytics (fatals + non-fatals + ANRs) + Analytics + Performance, classifies each finding by the §11.4.47 severity table, dedup-maps to existing Issues.md entries via the three-tier algorithm, and drafts new Issue entries for unrecognised findings. Skipping the pass is a §11.4 PASS-bluff: Firebase IS the captured evidence from real end-user devices; ignoring it is the precise “tests pass, feature broken in the wild” failure mode §11.4 prohibits.

Five mandatory elements (ALL must hold):

1. **Trigger cadence (5 trigger types).** Pre-build (blocking, <24 h freshness), pre-flash (blocking, <24 h), pre-tag (blocking, <6 h), daily 09:00 UTC default (non-blocking sweep), post-deployment burn-in T+24 h (non-blocking).
2. **Three-source query.** Crashlytics + Analytics + Performance — ALL three. Skipping one source is a §11.4 PASS-bluff (Firebase reports a regression on the skipped axis and we never see it).
3. **Issues.md output.** Every “problematic” finding MUST map to either (a) a new Issues.md entry, OR (b) a recognised existing entry via the dedup algorithm. Both paths produce a full Firebase Console URL in the entry’s metadata so any future agent can re-fetch the raw data without rebuilding context.
4. **Three-tier deduplication.** Tier 1 (exact Firebase Issue-ID match) → Tier 2 (stacktrace-similarity cluster hash: top-3 non-generic frames normalised + crash class → SHA-256 first 8 hex chars) → Tier 3 (monthly operator merge review). Multiple Issues.md entries pointing at the same underlying root cause is a §11.4 violation (operator chases ghosts; root cause stays unfixed).
5. **Comprehensive root-cause analysis.** Every Firebase-sourced Issue carries the §11.4.4(a) systematic-debugging output — Phase 1 evidence, Phase 2 pattern, Phase 3 root-cause hypothesis (UNCONFIRMED until §11.4.43 RED test reproduces), Phase 4 fix direction. Comment-only entries (“crash in App X”) are PASS-bluff stubs and FAIL the Issue-xref gate.

Severity classification: Crashlytics FATAL impactedUsers > 10 last 7d → Critical; 1–10 → Major; 0 but historic → Low. NON-FATAL > 100/day → Major; 10–100 → Low; < 10 → Informational. ANR > 1% sessions → Major; 0.1–1% → Low. Performance KPI regression > 20% → Major; 10–20% → Low. Analytics funnel-drop > 50% → Major; 25–50% → Low; feature_used regression > 75% → Major. Composes with §11.4.16 Type assignment.

Gates:

- CM-COVENANT-114-47-PROPAGATION — anchor present across every CLAUDE.md / AGENTS.md in the covenant file set.

- CM-AF-FIREBASE-REVIEW-CADENCE — scripts/firebase/review_round.sh exists, executable, contains the 7-stage pipeline literals + dedup algorithm + severity table.
- CM-AF-FIREBASE-ISSUE-XREF — every Issues.md entry whose ****Source:**** is Firebase Crashlytics / Firebase Analytics / Firebase Performance carries ****Firebase Issue IDs:**** + ****Firebase URL:**** + (****Stacktrace Cluster Hash:**** for Crashlytics OR ****KPI:**** for Performance OR ****Funnel:**** for Analytics).

Paired mutations strip the anchor / rename the review_round function / remove the Firebase-URL field from the template → respective gate FAILs.

Composition. §11.4.4 (STOP-on-discovery), §11.4.4(a) (systematic-debugging), §11.4.6 (no-guessing — Firebase data IS captured evidence), §11.4.7 (demotion-evidence — a previously- Fixed Issue resurfacing in Firebase IS positive captured evidence the fix didn't hold; reopen attribution = “AI: captured-evidence-contradicts” per §11.4.34), §11.4.10 (credentials — never log the bearer token), §11.4.12 (Issues_Summary sync), §11.4.14 (cleanup of /tmp/firebase_review_* work-dirs), §11.4.15 (Status: Queued), §11.4.16 (Type per severity), §11.4.34 (Reopened-Details on Firebase-resurface), §11.4.42 (implicit step 2 inclusion), §11.4.43 (RED test per stacktrace before fix lands), §11.4.44 (revision header on every Firebase-sourced Issue), §11.4.45 (Firebase Status.md at docs/firebase/status/Status.md), §11.4.46 (validation pass consults latest Firebase delta).

No escape hatch. No `--skip-firebase-review`, `--no-issue-from- firebase`, `--firebase-review-not-applicable` flag is permitted. The operator MAY filter by minimum severity (`--severity-min major`) to reduce noise but the pass itself MUST execute.

Classification: universal (per §11.4.17) — every consuming project that ships to real end-user devices benefits from the same review cadence and the same dedup algorithm.

§11.4.48 — UI-Driven Video Testing Mandate (User mandate, 2026-05-18)

Forensic anchor — verbatim user mandate (2026-05-18):

“We MUST make sure that all video playback testing with 2nd display is performed by fully automatically using the application's UI / UX, with full navigation, choice of video and clicks to real UI buttons to play or switch / stop videos being played! Maybe direct execution of Intents or Broadcasts does not have reported issues! We MUST test ALL WAYS users or Systems may / could play some video contents or streams! For all applications, all video and audio stream types! All supported codecs! EVERYTHING verified with 2nd display on D3 device and every other device with connected 2nd display and Avrus for proper codec(s) used in its Web Interface / Dashboard! Write as much as possible new tests for all supported test types!”

Why the existing video-test set is insufficient. The legacy `test_video_routing.sh` (v33), `test_video_secondary_display.sh`, `test_video_playback_automation.sh` (v1) plus the 53 `per_app_video` wrappers all dispatch playback via `am start -a android.intent.action.VIEW` or via `MediaSession cmd media_session play` — surfaces that skip the app's own

content-selection UI, player-surface allocation, MediaSession lifecycle (start/pause/stop), and back-button cleanup. A defect like §CL routing-race or §CM frozen-frame mid-back-press cannot reproduce via Intent dispatch because Intent dispatch re-launches the app cold each time and `am force-stop` is a clean cleanup path that masks the leak.

Operative rule. Every video-playback test that asserts secondary-display routing MUST traverse the user-equivalent UI path. Mandatory 5-surface coverage per video app:

1. **Surface A — Launch via launcher icon.** `am start -n <launcher-pkg>/<launcher-activity>` then `uiautomator dump` → locate app icon tile → `input tap X Y`. NOT `am start <video-pkg>/<MainActivity>` (skips launcher).
2. **Surface B — Navigate to content list.** Per-app driver script knows the app's home-tab → content-tab swipe / tap sequence. Asserts at least one content tile is visible via `uiautomator dump` content-desc / text match.
3. **Surface C — Select + play.** Tap a specific content tile (driver records the picked tile's content-desc as captured evidence). Wait for player surface visible.
4. **Surface D — In-app control interaction.** Pause via in-app pause button (not `input keyevent KEYCODE_MEDIA_PAUSE`), resume via in-app play button, in-app rapid switch to a second video (covers §CL).
5. **Surface E — Stop via back button.** `input keyevent KEYCODE_BACK` to exit player. Assert §11.4.14 cleanup: `VOM activeDecoder == null`, secondary surface cleared, no orphan MediaSession.

App-coverage matrix: every video-capable app in `PRODUCT_PACKAGES` MUST have a UI-driven driver script set (Layer 2 per Section 7 design) AND at least one scenario that uses it. Today: 15 apps covered fully in the initial batch (VK Video, MPV, Lampa as fully-implemented; remaining 12 as templated stubs per Section 8). Expansion to all 53 video-app entries is queued as follow-up work under §O Issues.md per-app matrix.

Stream-type matrix: progressive HTTP / HLS / DASH / RTMP / file- local / DRM-protected. Per app, the driver script tags which stream types its picked content exercises. Aggregator emits a coverage matrix per cycle so gaps are visible.

Codec matrix: H.264 / H.265 / VP9 / AV1 / MPEG-2 / MP4V (video) + AC-3 / E-AC-3 / TrueHD / DTS / DTS-HD / MLP / Opus / AAC (audio). Per scenario, the codec assertion is made via `dumpsys media .metrics` (decoder name in event log) AND via Arvus codec-state probe (audio side) per §11.4.13.

Secondary-display verification. On any device with HDMI-A-1 attached (detected via `dumpsys display | grep "Display 2"`): - Captured-evidence: `dual_display_record.sh` ON for both displays per §11.4.5 - Assertion: `ffprobe -count_frames` on secondary mp4 reports > N frames within 3 s of play-action timestamp - Anti-assertion: `ffprobe -count_frames` on primary mp4 reports ONLY launcher/home pixels during playback window (no video-decoder output) - VOM cross-check: `dumpsys video_output_manager` shows `activeDecoder != null` during playback and `== null` after stop within 3 s

When secondary absent (D2 default topology): SKIP per §11.4.3 with explicit reason “topology: no secondary display attached”. NEVER FAIL on missing topology.

Arvus codec-state mandate. For every test asserting an audio codec (AC-3, E-AC-3, TrueHD, DTS, etc.) — composed with §11.4.13 + §CG: - During playback window, call `arvus_probe_codec_state` ≥ 3 times at 1 s intervals (transients fade by sample 2-3) - Call `arvus_screenshot_capture` per §CG to attach visual evidence of the dashboard - Assert reported codec matches expected via `arvus_assert_codec_format <regex>` - ARVUS_HOST unreachable (ABK4 not joined / sink off) → SKIP per §11.4.3, never FAIL

Pre-build gates: - CM-COVENANT-114-48-PROPAGATION — asserts anchor literal present in every CLAUDE.md / AGENTS.md across parent + 10 owned submodules + HelixQA dependencies (42-file scan). - CM-AF-UI-DRIVEN-VIDEO-COVERAGE — asserts the directory `device/rockchip/rk3588/tests/ui_driven/` exists, contains `lib/ui_driver.sh` reference + `scenarios/` subdir + per-app driver subdirs for at least the 3 reference apps (`vk_video`, `mpv`, `lampa`), and that `scripts/testing/` `run_ui_driven_video_suite.sh` is executable. - Paired mutations: deleting `ui_driver.sh` → CM-AF-UI-DRIVEN- VIDEO-COVERAGE FAILs; stripping the anchor literal → CM-COVENANT-114-48-PROPAGATION FAILs.

No escape hatch. No `--use-intent-shortcut` / `--skip-ui-traverse` / `--legacy-intent-mode` flag exists. The discipline exists because Intent-based dispatch is exactly the class of test that PASSES while users hit real defects (§11.4 forensic anchor). Operators who feel UI-driven tests are “too slow” should land scenario-level parallelism (multi-device fan-out) rather than bypass the discipline.

Composition. - §11.4.3 — topology dispatch (secondary present vs absent) - §11.4.5 — captured-evidence quality (every UI tap captured) - §11.4.13 — Arvus codec-state mandate - §11.4.14 — playback cleanup verified via UI back-button path - §11.4.43 — RED-test-first discipline (initial UI tests RED against §CL/§CM) - §11.4.44 — revision header on every plan/driver doc - §CG — Arvus dashboard screenshot

Classification: universal (per §11.4.17). Applies to every project that consumes the constitution submodule AND ships any video-capable Android app.

§11.4.49 — Dual-Approach Testing Mandate (User mandate, 2026-05-18)

Forensic anchor — verbatim user mandate (2026-05-18):

“Kinopoisk playback of 5.1 audio supported movies MUST BE done via UI / UX full automation and via execution of Intents (or Broadcasts) directly both! We could have shared tests base and specialized part which will run / play the movie with properly chose audio track (5.1). Document everything up to the smallest details! Create as much as needed new tests for all supported test types and Challenges!”

Composition with §11.4.48. §11.4.48 mandated UI-driven traversal for every video routing test, eliminating the Intent-only PASS-bluffs. §11.4.49 REFINES rather than replaces: every feature test ships in BOTH variants — UI-driven AND Intent-driven — over a shared assertion base. Either alone is a §11.4 PASS-bluff for the OPPOSITE half of the stack. UI catches app-side bugs (content selection, player surface allocation, MediaSession lifecycle, in-app overlays, back-button cleanup).

Intent catches framework/system-server bugs (Intent extras parsing, MediaCodec.configure hook MIME routing, IVideoOutputManager bind timing, broadcast permission gating, headless cron automation paths).

Operative rule — 5 mandatory elements:

1. **Both variants required.** Every feature test exercising a user-visible behaviour MUST ship both `<feature>_ui.sh` (uiautomator-based, §11.4.48 surfaces A–E) AND `test_<feature>_intent.sh` (am start --es / am broadcast-based). Either alone is forbidden.
2. **Shared assertion base.** Codec-state assertions, captured- evidence collection (screen-recording, ffprobe, RMS amplitude analysis), Arvus codec-state probe + dashboard screenshot, and §11.4.14 cleanup MUST be implemented in a single POSIX-sh library (`tests/lib/dual_approach_test_base.sh`) and reused by BOTH variants of every dual-approach test. Code duplication in the shared layer is forbidden.
3. **Specialised driver code.** UI variant uses §11.4.48 per-app driver scripts under `tests/ui_driven/<app>/`. Intent variant uses `am start --es content_id <id> --es audio_track_id <track> / am broadcast -a <action>` directly. Each variant’s specialised layer is responsible ONLY for “how the playback gets started”; everything downstream of “playback is now happening” routes through the shared base.
4. **Comprehensive documentation per test.** Every dual-approach test set ships with: (a) §11.4.44 revision header in BOTH variant scripts, (b) per-feature contract under `docs/dual_approach/ <feature>.md` describing the captured-evidence pairing and operator-side verification, (c) `Issues.md / Fixed.md` entry cross-linking both variant paths.
5. **Kinopoisk 5.1 EAC3 is the canonical first implementation.** The shared base + Kinopoisk 5.1 UI variant + Kinopoisk 5.1 Intent variant land in the same batch as this mandate. Subsequent features (Netflix Dolby Atmos, MPV DTS-HD, Lampa+TorrServe HEVC, VLC FLAC) port to the same pattern. The §CN subagent’s fix to the underlying decoder pipeline is the FIRST GREEN target of these tests; both variants are RED per §11.4.43 until §CN lands.

Captured-evidence directory contract. Both variants write to mirror-structured directories `qa-results/dual_approach/<F>/<run- ts>/{ui,intent}/`. Identical filenames; orchestrator diffs the two `result.json` files. Status mismatch (UI=PASS but Intent=FAIL or vice-versa) is itself a finding — it pinpoints which half of the stack contains the bug.

Pre-build gates: - CM-COVENANT-114-49-PROPAGATION — anchor literal in every CLAUDE.md / AGENTS.md across parent + 10 owned submodules + HelixQA dependencies (42-file scan). - CM-AF-DUAL-APPROACH-COVERAGE — `tests/lib/dual_approach_test_base.sh` exists + sources `anti_bluff/ui_driver/credentials` + exports `dat_init / dat_start_capture / dat_assert_codec_state / dat_assert_video_frames / dat_assert_audio_channels / dat_arvus_dashboard_capture / dat_cleanup / dat_report_finding`. - CM-AF-KINOP0ISK-5-1-DUAL-COVERAGE — both `tests/ui_driven/kinopoisk/kinopoisk_5_1_play_movie_ui.sh` AND `tests/test_kinopoisk_5_1_play_movie_intent.sh` exist + both source the shared base + both reference EAC3 codec + both assert 5.1 / 6-channel. - Paired mutations (3): strip anchor → propagation gate FAILS; delete UI variant → coverage gate FAILS; delete Intent variant → coverage gate FAILS.

No escape hatch. No `--ui-only` / `--intent-only` / `--skip-` dual flag exists. The discipline exists because either-alone is the PASS-bluff pattern §11.4 specifically prohibits. Operators who feel dual variants are “too slow” should land scenario-level parallelism rather than bypass the discipline.

Composition. - §11.4.3 — topology dispatch (Arvus reachable, secondary present) - §11.4.5 — captured-evidence content-quality analysis - §11.4.6 — UNCONFIRMED tagging for un-verified Kinopoisk element IDs - §11.4.10 — credentials never echoed, per-device per-service loader - §11.4.13 — Arvus codec-state mandate (out-of-band evidence) - §11.4.14 — playback cleanup via EXIT trap in shared base - §11.4.17 — universal classification - §11.4.43 — RED-first TDD (both variants RED until §CN fix lands) - §11.4.44 — revision header on every variant script + design doc - §11.4.48 — UI-driven traversal (this MANDATE refines, not replaces) - §CG — Arvus dashboard screenshot via Presenter receiver - §CB — credentials loader

Classification: universal (per §11.4.17). Applies to every project that consumes the constitution submodule AND ships any Android app whose user-visible behaviour can be triggered through both UI and Intent/Broadcast paths.

§11.4.50 — Deterministic Consistency Mandate (User mandate, 2026-05-18)

Forensic anchor — verbatim user mandate (2026-05-18):

“Our anti-bluff / bluff-proofed / proof-driven work **MUST** satisfy the following: not a single feature, System or application flow, use case, edge case or procedural action(s) **MUST NOT** partially work, or sometimes work and sometimes not! There is only one truth that **MUST BE** fulfilled - no matter how many times we repeat all these scenarios with variations in data, System state(s) and speed of execution, results **MUST BE** consistent and successful without exception! No false positives or bluffing of any kind is allowed! We **MUST** add as much full automation tests which will use real System and Applications UI and UX with all flows, use cases and edge cases as needed or as much as it is possible!”

Why this anchor exists. §11.4.7 already forbids the vocabulary intermittent / transient / flake in closure narratives, but that’s a textual ban — operators could still let a test “PASS once, FAIL once, PASS again” and report only the first PASS. §11.4.50 closes that gap **MECHANICALLY** by requiring every PASS to come from N identical iterations rather than from a single observation.

Operative rule — 5 mandatory elements:

1. **N-iteration deterministic outcome.** Every test that **PASSes** **MUST** have been executed N times (default N=3, N=10 for cycle- validation suites) against the same firmware MD5 + same device
 - same topology and produced **IDENTICAL** PASS in every iteration. A test whose outcome diverges across iterations is auto-FAIL per this mandate — there is no “first **PASSed** therefore X was a flake” path. §11.4.7’s expanded forbidden-vocabulary list is enforced mechanically here.

2. **Edge-case + flow + use-case coverage.** Every public API path (Activity / Service / Broadcast receiver / ContentProvider URI / IPC interface / JNI entry / sysprop write / sysfs node / init.rc trigger) MUST have ≥ 1 dedicated test that drives it. Untested paths surface in a feature-coverage-matrix audit and block release at the 99% threshold (ratchet sequence $70 \rightarrow 85 \rightarrow 95 \rightarrow 99$ mirrors §11.4.18 the existing project-side docs-coverage ratchet).
3. **Reliability check helper.** Project anti-bluff helper library ships `ab_run_n_times <test_name> <N> <fn> [args...]`:
 - Loops N times, captures exit code + evidence-hash per iter
 - Asserts all N exit codes identical AND all N evidence-hashes identical
 - On any divergence: `ab_fail` with full N-iteration report
 - Per §11.4.7 forbidden vocabulary: NO operator-facing escape path converts a divergent N-iter run into PASS
4. **Feature-coverage-matrix audit.** Project ships `scripts/testing/feature_coverage_audit.sh`:
 - Walks source layers (every public API surface)
 - Walks test layers (`test_*.sh` references)
 - Emits coverage report at `qa-results/feature_coverage/ <ISO-UTC>/coverage.tsv`
 - Pre-build gate enforces minimum threshold; threshold ratchets up each phase ($70 \rightarrow 85 \rightarrow 95 \rightarrow 99$)
5. **Composes with every prior anti-bluff anchor.** §11.4.50 does NOT replace §11.4.43 / §11.4.48 / §11.4.49 — it adds the N-iter consistency dimension across all of them. RED-first TDD (§11.4.43) still applies: every new test starts RED and becomes GREEN after the fix lands. UI-driven (§11.4.48) and dual-approach (§11.4.49) still apply: every test ships in both variants AND each variant passes N iterations identically.

Pre-build gates:

- CM-COVENANT-114-50-PROPAGATION — anchor literal §11.4.50 present in every CLAUDE.md / AGENTS.md across parent + 10 owned submodules + HelixQA dependencies (≥ 42 -file scan; phase 1 uses 5-canonical-file scan matching §11.4.49 propagation pattern).
- CM-AF-RELIABILITY-CHECK-WIRED — `ab_run_n_times` function literal present in `device/rockchip/rk3588/tests/lib/anti_bluff.sh` AND at least 3 on-device tests source-reference it. The 3-test floor ratchets upward each phase.
- CM-AF-FEATURE-COVERAGE-MATRIX — `scripts/testing/feature_coverage_audit.sh` exists + executable + has run within the last 7 days + most-recent coverage report meets the current threshold.

Paired mutations (3):

- Strip the §11.4.50 anchor literal from `constitution/CLAUDE.md` via `sed -i 's|11.4.50|11.4.MUTATED|g' → CM-COVENANT-114-50-PROPAGATION MUST FAIL.`
- Rename `ab_run_n_times` to `ab_run_n_times_DISABLED` in `anti_bluff.sh` via targeted `sed → CM-AF-RELIABILITY-CHECK- WIRED MUST FAIL.`

- Move scripts/testing/feature_coverage_audit.sh aside via mv → CM-AF-FEATURE-COVERAGE-MATRIX MUST FAIL.

No escape hatch. No --allow-flake, --first-pass-suffices, --skip-n-iter, --skip-coverage-audit flag exists. The discipline exists because the user mandate is unambiguous: “results MUST BE consistent and successful without exception.”

Composition. - §11.4.1 — FAIL-bluffs forbidden - §11.4.2 — captured-evidence (every iteration captures its own) - §11.4.5 — quality analysis (per-iteration content compared) - §11.4.6 — no guessing (UNCONFIRMED required for any iteration-budget assumption) - §11.4.7 — forbidden flake/intermittent vocabulary (enforced mechanically by this mandate) - §11.4.43 — RED-first TDD (every new test starts RED, N-iter applies across the GREEN transition) - §11.4.46 — validate-before-suite (N-iter applied at single-test scope first) - §11.4.48 — UI-driven traversal (each UI traversal MUST be N-iter-consistent) - §11.4.49 — dual-approach (both variants of every dual-approach test MUST pass N iterations identically)

Classification: universal (per §11.4.17). Applies to every project that consumes the constitution submodule.

§11.4.51 — Live-ADB-First Maximization Mandate (User mandate, 2026-05-18)

Forensic anchor — verbatim user mandate (2026-05-18):

“Was it possible to max the fix and test it via ADB or it can be only done by making the fix and reflashing for validation and verification? If such option exists, we should always make the best of it! Making fix on live devices via ADB → Once done commit and push changes, rebuild System and reflash → Validate and verify everything (with the test we write as well). Make sure this idea(s) is (are) part of our root (constitution Submodule) Constitution, AGENTS.MD and CLAUDE.MD if they are not already! EVERY DETAIL IS IMPORTANT!!!! Not a single idea or idea’s detail can be ignored, skipped, relativized or bluffed!”

§11.4.51 REFINES §11.4.43 step 2 (“LIVE-ADB-PROBE — try the fix on the running device first”) with mechanical enforcement: a per-file- class decision matrix, a classifier helper, and a commit-message footer literal.

Operative rule (5 mandatory elements):

1. **Classify every fix** by rebuild-requirement using the project’s per-file-class decision matrix. No guessing (§11.4.6). Every file in the diff cites the matrix row that classified it.
2. **If LIVE_ADB_TESTABLE**, the operator MUST apply the fix to the running device via the appropriate adb channel first (adb push for scripts, setprop persist.* for runtime properties, mount -o remount, rw /vendor for boot scripts, pm install -r for APKs built locally via gradle). Run the §11.4.43 RED test against the live-probed device, capture positive-evidence PASS,

THEN commit source-side + rebuild + reflash as belt-and-suspenders re-validation. Commit message footer MUST state `LIVE_ADB_VALIDATED: yes` with one-line description of the live-probe sequence.

3. **If `REQUIRES_REBUILD`**, the operator proceeds directly to source-side fix + rebuild + flash. Commit message footer MUST state `REQUIRES_REBUILD: <reason>` citing the matrix row that classified the file. Categories of `REQUIRES_REBUILD`: kernel, AOSP framework Java / AIDL, native C++ inside APEX, sepolicy, init.rc, `ro.*` build properties (immutable post-boot), `Android.bp` / `Android.mk` / `BoardConfig.mk`, XML resource overlays, codec XML inside APEX, ramdisk-bundled artifacts.
4. **If `mixed batch`**, commit footer states `LIVE_ADB_VALIDATED: partial` and enumerates per-file classification. Per §11.4.9 batching is preferred; live-probe the testable files first, then batch with the rebuild-required ones into a single rebuild.
5. **Helper script enforces classification mechanically.** The project provides `classify_fix_rebuild_requirement.sh` which walks `git diff --name-only`, looks up each file against the matrix, and emits the recommended commit-message footer. No operator-facing override flag converts an unclassified path to `LIVE_ADB_TESTABLE` by default — unmatched paths classify as `REQUIRES_REBUILD: unmatched-path` (safe default per §11.4.6).

Per-file-class decision matrix (canonical):

File class	Rebuild required?	Reason
<code>*test_*.sh</code> / <code>tests/lib/*.sh</code> (on-device)	NO	adb push to <code>/data/local/tmp/tests/</code>
Host-side test scripts	NO	host script — no device interaction
<code>atmosphere-*.sh</code> boot scripts	DEPENDS	adb push to <code>/vendor/bin/</code> after remount
Markdown docs / Constitution / plans / Issues / Fixed / CONTINUATION	NO	host docs — no firmware impact
Forked-player Kotlin (Presenter, MPV, SmartTube, TorrServe, Lampa, VLC, etc.)	YES via <code>gradle local build + adb install -r</code>	APK re-link required but no full firmware rebuild
Framework Java (frameworks/base/**/*. <code>java</code>)	YES	<code>system_server</code> requires reboot at minimum
AIDL (<code>*.aidl</code>)	YES	stub generation requires recompile
Native C++ (external/**/*. <code>cpp</code> , frameworks/native/**/*. <code>cpp</code>)	YES	builds into APEX or system library
	YES	squashfs read-only by Mainline integrity

File class	Rebuild required?	Reason
APEX libraries (vendor/ **/lib/**/*.*.so inside APEX path)		
Kernel sources (kernel-5.10/**)	YES	requires kernel rebuild + boot.img
sepolicy (*.te)	YES	requires policy build
init.rc	YES	parsed at boot only
ro.* build properties	YES	immutable post-boot
persist.* runtime properties	NO	setprop persist.* mutable
XML resource overlays	YES	RRO recompile
media_codecs_*.xml	YES	APEX-bundled
display_settings.xml (userdata)	DEPENDS	adb push possible if SELinux permits
Android.bp / Android.mk / device.mk / BoardConfig.mk	YES	rebuilt-into-image artifacts
Test fixture binary assets	NO	adb push to /data/local/ tmp/

Pre-build gates: CM-COVENANT-114-51-PROPAGATION (anchor across canonical files) + CM-AF-CLASSIFY-FIX-HELPER-EXISTS (helper present + sentinel literals) + CM-AF-LIVE-ADB-FIRST-COMMIT-MARKER (advisory WARN that scans recent commits for the footer literal). Three paired meta-test mutations.

Composition. - §11.4.43 — TDD-fix workflow (§11.4.51 REFINES step 2 with the mechanical classifier). - §11.4.9 — batch-source-fixes-before-rebuild (compose: live-ADB- test individually, then batch-commit). - §11.4.6 — no guessing (each classification justified by matrix row; unmatched paths default to rebuild-required, never silent testable). - §11.4.46 — validate-recent-work-before-post-flash (live-probe IS implicit pre-validation). - §11.4.48 — UI-driven tests (live-probe CAN use uiautomator over adb shell). - §11.4.49 — dual-approach (live-probe applies to both UI and Intent variants). - §11.4.50 — deterministic consistency (live-probe runs N iterations on the live device before commit).

No escape hatch — no --skip-classify, --assume-rebuild, --no-footer-required flag exists. The discipline exists because the user mandate is unambiguous: “EVERY DETAIL IS IMPORTANT!!!! Not a single idea or idea’s detail can be ignored, skipped, relativized or bluffed!”

Classification: universal (per §11.4.17). Applies to every project that consumes the constitution submodule — the matrix itself MAY be project-specific (each consumer adapts the file-class list to its own source tree), but the mandate to classify-before-commit + the LIVE_ADB_VALIDATED / REQUIRES_REBUILD footer literals + the classify-helper-script + the three pre-build gates are universal.

§11.4.52 — Autonomous-Validation Mandate (User mandate, 2026-05-18)

Forensic anchor — verbatim user mandate (2026-05-18):

“Make sure we have full automation tests which will do all this work in full automation! IMPORTANT: Make sure that all existing tests and Challenges do work in anti-bluff manner — they MUST confirm that all tested codebase really works as expected! We had been in position that all tests do execute with success and all Challenges as well, but in reality the most of the features does not work and can’t be used! This MUST NOT be the case and execution of tests and Challenges MUST guarantee the quality, the completion and full usability by end users of the product! This MUST BE part of Constitution of our project, its CLAUDE.MD and AGENTS.MD if it is not there already, and to be applied to all Submodules’s Constitution, CLAUDE.MD and AGENTS.MD as well.”

Why this anchor exists. §11.4.25 (full-automation coverage) and §11.4.27 (no-fakes-beyond-unit-tests + 100%-test-type-coverage) and §11.4.39 (per-feature on-device end-user validation) collectively mandated that every user-facing feature has automation tests with captured runtime evidence. They did NOT explicitly forbid operator-attended-only validation paths — i.e., a feature whose ONLY proof of working state requires a human to physically present at the device, drive UI manually, and observe outcomes. Phase 39.§CN exposed the gap: §CN’s Kinopoisk 5.1 EAC3 closure was structurally fixed (libavutil bundled in APEX, six decoders register) yet end-user-validation degraded to “operator drives remote control while test polls” because Kinopoisk’s TV-Compose UI is not accessibility-instrumented and uiautomator returns near-empty hierarchy. That validation path is operator-attended, non-reproducible in CI, and cannot prove deterministic consistency (§11.4.50) across N iterations. The User mandate elevates “full automation” from “preferred” to “mandatory”: every user-facing PASS MUST have at least one captured-evidence path that does NOT require operator presence.

Operative rule (4 mandatory elements):

1. **Every user-facing feature MUST have at least one autonomous validation path.** “Autonomous” means: the validation runs end-to-end via `adb shell` + scripted automation, produces captured runtime evidence per §11.4.5, and reaches a PASS / FAIL verdict WITHOUT a human present to drive UI, observe screen, or make decisions. Operator-attended tests are SUPPLEMENTARY, never PRIMARY. A feature whose ONLY validation path is operator-attended is a §11.4.52 violation regardless of how carefully the operator’s captures match §11.4.5 evidence requirements — the path does not scale to CI, does not run on every commit, does not survive operator unavailability, and produces the exact “tests pass but feature doesn’t work for users” failure mode §11.4 specifically forbids.
2. **Acceptable autonomous-validation paths** (any one or more — non-exhaustive):
 - **Programmatic instrumentation APK** — a small platform-signed app that exercises the under-test surface directly via SDK APIs (e.g., `MediaCodec.createDecoderByName`, `MediaCodec.configure`, `AudioTrack` writes) and writes a structured JSON result file to a

discoverable path. Closes UI-instrumentation gaps where the target app is not accessibility-instrumented.

- **Headless intent dispatch + state poll** — `am start --es / am broadcast --es` to dispatch the under-test code path, then poll `dumpsys, /proc/<pid>/maps, media.metrics, /sys/...`, kernel sysfs, or the integration's network probe for the expected state transition.
- **ADB-driven uiautomator** — applicable ONLY if the target app exposes accessibility nodes the dump can resolve. MUST verify via `uiautomator dump | grep -c clickable=true ≥ 1` before claiming uiautomator coverage is real. A near-empty hierarchy is itself the captured-evidence proving UI-driven automation is INFEASIBLE for that target and the test MUST fall back to instrumentation-APK or headless-intent paths.
- **Network-side sink probe** — for features that surface on a network-accessible peripheral (e.g., Arvus HDMI dashboard at `http://<sink>/`, Sonos REST API, AirPlay receiver), the sink's own report counts as autonomous out-of-band evidence per §11.4.13.
- **HelixQA autonomous QA session** — for projects that have HelixQA fully incorporated per §11.4.27, the orchestrator's end-to-end session output counts as autonomous evidence.

3. **Per-feature classification + tracking.** Every entry in the project's coverage ledger (§11.4.25) MUST classify each feature's autonomous-validation path as one of: `AUTONOMOUS_VERIFIED` (path exists + last-run-green per §11.4.46), `AUTONOMOUS_DESIGNED` (path exists but RED per §11.4.43), `OPERATOR_ATTENDED_ONLY` (path requires human presence — release blocker until promoted to one of the other states), or `NOT_APPLICABLE` (e.g., a CLI tool that has no UI surface and the unit + integration tests already constitute autonomous proof). `OPERATOR_ATTENDED_ONLY` rows MUST cite a tracked work item per §11.4.15 + §11.4.16 that designs the autonomous-path migration. Closing the work item requires the migration to land, not just the operator capture.

4. **Anti-bluff for autonomous paths themselves.** An autonomous path that returns PASS without exercising the user-visible behaviour is a §11.4.52 violation just as severe as the original §11.4 PASS-bluff pattern. The autonomous path MUST produce positive captured evidence per §11.4.5 (audio: channel count + sample rate + glitch census; video: frame count + routing target

- frame health + obstruction census). A grep on a metadata field that doesn't prove behaviour is not an autonomous validation — it is a metadata bluff dressed up as automation. Paired meta-test mutations (§1.1) MUST catch the autonomous path's degradation to metadata-only or grep-only.

Composition. - §11.4.25 — full-automation-coverage mandate. §11.4.52 is the strict expansion of invariant 1 + 2: “anti-bluff posture” and “proof of working capability” cannot be carried by operator-attended-only evidence — they MUST be carried by at least one autonomous path. - §11.4.27 — no-fakes-beyond-unit-tests + 100%-test-type-coverage. §11.4.52 is the operational layer that closes the gap between “the project has full-automation tests” and “every PASS has an autonomous evidence path”. - §11.4.39 — per-feature on-device end-user validation. §11.4.52 refines: the on-device scenario MUST itself be autonomously drivable. - §11.4.43 — TDD-fix-discipline. The autonomous path MUST exist in the RED state BEFORE the fix lands, then go GREEN after. An autonomous path authored after the fix is a §11.4 PASS-bluff. - §11.4.48 — UI-driven video testing. §11.4.52 establishes that when uiautomator returns a near-empty hierarchy on the target app, UI-driven IS

infeasible and the test MUST fall back to instrumentation-APK or headless-intent. The fallback is REQUIRED — not optional. - §11.4.49 — dual-approach testing. The Intent variant in a dual-approach pair often IS the autonomous path when the UI variant requires operator assistance. - §11.4.50 — deterministic consistency. Autonomous paths run N iterations naturally; operator-attended paths cannot satisfy the N-iteration discipline at scale. - §11.4.51 — live-ADB-first maximization. The instrumentation-APK path is itself LIVE_ADB_TESTABLE (gradle build + adb install -r + run + read JSON result), so live-probe IS the design loop for the autonomous path.

Pre-build gates: - CM-COVENANT-114-52-PROPAGATION — anchor literal §11.4.52 present in all 5 canonical files (constitution/Constitution.md, constitution/CLAUDE.md, constitution/AGENTS.md, plus the parent consumer's CLAUDE.md + AGENTS.md) PLUS the full per-consumer propagation (parent project + every owned submodule + nested submodules + HelixQA dependencies — same scan pattern as CM-COVENANT-114-51-PROPAGATION). - CM-AF-AUTONOMOUS-PATH-PER-FEATURE — every entry in the consuming project's coverage ledger (§11.4.25) has a non-empty autonomous-classification column with value in {AUTONOMOUS_VERIFIED, AUTONOMOUS_DESIGNED, OPERATOR_ATTENDED_ONLY, NOT_APPLICABLE}. OPERATOR_ATTENDED_ONLY rows MUST cite a tracked migration work item. (Project-side gate — consuming projects implement the ledger surface in their own pre-build test suite.)

Paired mutations (per §1.1): - Strip §11.4.52 literal from constitution/Constitution.md → CM-COVENANT-114-52-PROPAGATION FAILs. - Inject a feature row with OPERATOR_ATTENDED_ONLY and NO tracked migration work item → CM-AF-AUTONOMOUS-PATH-PER-FEATURE FAILs.

No escape hatch. No --allow-operator-attended-only, --skip-autonomous-path, --manual-validation-suffices flag exists. The discipline exists because the user mandate is unambiguous: “execution of tests and Challenges MUST guarantee the quality, the completion and full usability by end users”. Operator-attended-only validation cannot guarantee that property at the release-gate layer — only autonomous paths can.

Classification: universal (per §11.4.17). Applies to every project consuming the constitution submodule. Each consumer adapts the autonomous-validation path catalogue to its own technology surface (Android: instrumentation APK + uiautomator + headless intent; Web: Playwright headless + API probe; CLI: subprocess + stdout/stderr assertion; mobile native: device farm headless runner) but the mandate to provide at least one autonomous path per user-facing feature is universal.

§11.4.53 — Fixed_Summary parity mandate (User mandate, 2026-05-18)

Forensic anchor — verbatim user mandate (2026-05-18T17:55Z):

“Note: Just like for Issues we have Issues_Summary, for Fixed we MUST HAVE Fixed_Summary - like all other docs: ALWAYS in sync and up to date and ALWAYS exported into the PDF and HTML! Add this mandatory rule / constraint into the root (constitution Submodule) Constitution, AGENTS.MD and CLAUDE.MD.”

Why this anchor exists. §11.4.12 already established that docs/Issues_Summary.md is the canonical short-form summary of docs/Issues.md and MUST always be regenerated + re-exported (HTML + PDF) whenever the source changes. §11.4.19 added the column- alignment requirement so Fixed_Summary.md mirrors Issues_Summary.md shape. §11.4.53 closes the propagation gap: the parity discipline that §11.4.12 imposes on Issues / Issues_Summary MUST apply symmetrically to Fixed / Fixed_Summary. The mechanical behaviour already exists in scripts/testing/sync_issues_docs.sh (lines 98- 105 regenerate Fixed_Summary in stage 1b alongside Issues_Summary in stage 1a), but the mandate was never explicit in the Constitution — and an unstated rule is a §11.4 PASS-bluff risk: future refactors of sync_issues_docs.sh could drop the Fixed_Summary half without violating any canonical authority. §11.4.53 makes the existing behaviour explicit canonical authority so the gate set can defend it.

Operative rule. docs/Fixed_Summary.md is the symmetric short- form summary of docs/Fixed.md. It MUST be regenerated whenever Fixed.md changes. Its HTML + PDF exports MUST travel with the markdown (identical mtimes within sync_issues_docs.sh granularity). Stale exports are §11.4.53 violations regardless of whether the underlying .md is correct — an operator (or future agent) reading the HTML or PDF gets a divergent view of which items are closed, and the §12.10 CONTINUATION resumption guarantee silently breaks. Same discipline as §11.4.12 Issues_Summary applied to Fixed.md.

Generator. scripts/testing/generate_fixed_summary.sh is the canonical generator. It MUST exist + be executable + emit a markdown table whose header columns include Status and Type (per §11.4.19 column-alignment) so the Fixed_Summary.md shape mirrors Issues_Summary.md exactly.

Auto-sync wrapper. scripts/testing/sync_issues_docs.sh is the single operator-facing entry point. It already regenerates BOTH Issues_Summary AND Fixed_Summary in one shot (stages 1a + 1b), then exports HTML + PDF (stage 2), then colorizes per §11.4.23 (stage 3), then re-renders the PDFs from the colorized HTML (stage 3b). MUST be invoked after any edit to Fixed.md (just as §11.4.12 requires after any edit to Issues.md). Manual invocation of just the Issues half (-- issues-only-style flag) is FORBIDDEN — the wrapper has no such flag and §11.4.53 prohibits adding one.

Sort order. Same pattern as §11.4.12 — by closure date DESC (most-recent-Fixed first), with §-letter / Fix-# secondary sort. Documented at the top of the generated file.

HTML + PDF travel-together. Per §11.4.12 + §11.4.44, the three file types (.md, .html, .pdf) MUST always have identical mtimes within sync_issues_docs.sh granularity. Stale exports violate §11.4.53 regardless of whether the underlying .md is correct.

Composition: - §11.4.12 — Issues_Summary parity is the sibling rule §11.4.53 symmetrizes. §11.4.12 and §11.4.53 are the canonical pair: edits to one source-of-truth (Issues.md or Fixed.md) trigger regeneration of BOTH summaries via the same wrapper. - §11.4.19 — atomic Issues→Fixed migration triggers Fixed_Summary regeneration. When an item closes (status Fixed (→ Fixed.md) / Implemented (→ Fixed.md) / Completed (→ Fixed.md) per §11.4.33), the operator moves the entry from Issues.md to Fixed.md atomically and runs sync_issues_docs.sh — both summaries refresh. - §11.4.23 — visual-cue & grouping colorizer post-processes both Issues_Summary.html AND

Fixed_Summary.html. Stale Fixed_Summary.html would carry pre-§11.4.23 styling — itself a §11.4.53 violation. - §11.4.33 — type-aware closure-status vocabulary. Fixed_Summary MUST respect the three terminal values (Fixed (→ Fixed.md), Implemented (→ Fixed.md), Completed (→ Fixed.md)) and emit them literally in the Status column. - §11.4.44 — revision-header mandate applies to Fixed_Summary.md exactly as it applies to every other tracked Markdown doc. - §12.10 — CONTINUATION.md resumption guarantee depends on the divergent-summary problem NOT existing; §11.4.53 is the explicit symmetric closure of that risk for the Fixed half.

Pre-build gates:

- CM-FIXED-SUMMARY-SYNC — canonical artifact-level gate (6 invariants): (1) docs/Fixed_Summary.md exists; (2) docs/Fixed_Summary.html exists with mtime ≥ Fixed_Summary.md mtime; (3) docs/Fixed_Summary.pdf exists with mtime ≥ Fixed_Summary.md mtime; (4) Fixed_Summary.md mtime ≥ Fixed.md mtime; (5) scripts/testing/generate_fixed_summary.sh exists + executable; (6) scripts/testing/sync_issues_docs.sh invokes the generator. Pre-build FAILs if any invariant is violated. Partially overlaps with CM-FIXED-COLUMN-ALIGNMENT (§11.4.19 Phase 39.AV gate) and CM-DOCS-EXPORT-SYNC (§11.4.15 4-doc + 6-export scan), but is explicitly the Fixed_Summary derived-doc parity invariant for §11.4.53 enforcement.
- CM-COVENANT-114-53-PROPAGATION — anchor literal §11.4.53 present across all 5 canonical files (constitution/Constitution.md, constitution/CLAUDE.md, constitution/AGENTS.md, parent consumer's CLAUDE.md + AGENTS.md). Same propagation pattern as CM-COVENANT-114-51-PROPAGATION and CM-COVENANT-114-52-PROPAGATION. Consumer projects propagate further (owned submodules, nested submodules, HelixQA deps) at their own per-project gate granularity.

Paired mutations (per §1.1):

- Strip §11.4.53 literal from constitution/CLAUDE.md → CM-COVENANT-114-53-PROPAGATION FAILs.
- Move scripts/testing/generate_fixed_summary.sh aside (so the invariant-(5) executable-check fails) → CM-FIXED-SUMMARY-SYNC FAILs.
- Touch docs/Fixed_Summary.md to a date older than docs/Fixed.md → invariant-(4) violated → CM-FIXED-SUMMARY-SYNC FAILs.

No escape hatch. No --skip-fixed-summary-sync, --issues-only, --summary-not-applicable flag exists. The discipline exists because the user mandate is unambiguous: “like all other docs: ALWAYS in sync and up to date and ALWAYS exported into the PDF and HTML”. A divergent Fixed_Summary breaks operator + AI-agent ability to discover what is fixed vs what is still open — and that breakage IS the §11.4 “tests pass but reality differs” pattern at the documentation layer.

Classification: universal (per §11.4.17). Applies to every project consuming the constitution submodule that maintains an Issues.md / Fixed.md split (the Issues-tracking lifecycle is already a §11.4.15 universal mandate, so this clause inherits the same universal scope). Projects that do not maintain a Fixed.md (e.g., single-binary throwaway prototypes) are NOT APPLICABLE; all projects that DO maintain a Fixed.md MUST land both gates and mutations within one cycle of adopting this Constitution version.

§11.4.54 — ATM-NNN ticket identifier mandate (User mandate, 2026-05-19)

Forensic anchor — verbatim user mandate (2026-05-19):

“Every workable item in Issues.md / Issues_Summary.md / Fixed.md / Fixed_Summary.md MUST carry a stable, unique, auto-incremental ATM-NNN ticket identifier. ATM- prefix, monotonic, never renumbered, append-only — once assigned to an item it stays bound to that item for the lifetime of the project across every reopen, migration, or rename.”

Why this anchor exists. §11.4.15 + §11.4.16 + §11.4.19 + §11.4.33 established the lifecycle vocabulary for tracked items, but the items themselves were identified only by §-letter or Fix-# — both of which are **unstable across migrations**. §-letter is heading-local; when an item moves from Issues.md to Fixed.md the §-letter may be reused for a different item. Fix-# is only assigned on closure for items that touched source code; tasks / features / pure-doc items have no stable identifier at all. Cross-references between commits, test logs, captured-evidence artifacts, Reopens.md history (§11.4.55), README.md doc-link rows (§11.4.57), and external systems (Firebase Crashlytics Issue-IDs, Atlassian-style tracker rows) all need an identifier that is **monotonic, stable, and project-wide unique**. §11.4.54 closes that gap.

Operative rule. Every workable item in docs/Issues.md AND docs/Fixed.md MUST carry a [ATM-NNN] ticket identifier in its heading, in the form ## §X.Y. [ATM-NNN] <title> (or ### §X.Y. for nested items). NNN is a positive integer, zero-padded to at least 3 digits (ATM-001, ATM-042, ATM-127, ...). Identifiers are allocated by a **canonical helper script** (scripts/testing/assign_atm_ticket_ids.sh) that maintains an append-only state file (scripts/testing/.atm_ticket_state.json). Once assigned to an item, an ATM-NNN MUST NEVER be:

1. **Renumbered** — the bind is permanent.
2. **Reused** for a different item even if the original item is force-deleted (the helper’s state file detects collisions).
3. **Decrement**ed — counter is monotonic-only.
4. **Skipped intentionally** — gaps are forbidden in the assignment sequence (gaps indicate state-file corruption and trigger gate FAIL).

State file format. scripts/testing/.atm_ticket_state.json is a JSON-lines file with one record per allocated ID:

```
{"atm_id": "ATM-001", "heading_hash": "<sha256 of normalized heading>",  
  "type": "Bug|Feature|Task", "current_location": "Issues|Fixed",  
  "current_status": "<lifecycle value>", "reopens_count": <int>,  
  "created_at": "<ISO 8601 UTC>", "last_modified": "<ISO 8601 UTC>"}
```

The heading_hash is the canonical binding key — when the helper scans Issues.md / Fixed.md, it computes the normalized hash of each heading and looks up an existing ATM-NNN before allocating a new one. This protects against accidental

renumbering on heading reflows / wording edits (the hash MUST stay stable; if the heading text changes substantially, the state file's `heading_hash` field is updated in place via an explicit migration command, NOT auto-rebound).

Summary-table column. `docs/Issues_Summary.md` and `docs/Fixed_Summary.md` MUST carry a leftmost ATM ID column showing the identifier. Generators (`generate_issues_summary.sh`, `generate_fixed_summary.sh`) MUST emit this column as the first data column so operators / agents can sort + filter on it.

Composition:

- §11.4.15 (Status), §11.4.16 (Type), §11.4.19 (column-alignment Issues/Fixed), §11.4.33 (type-aware closure terminal values). All four pre-existing mandates compose with §11.4.54 — the ATM-NNN is the new universal join key.
- §11.4.12 + §11.4.53 (Issues_Summary + Fixed_Summary regen on source changes) — the helper MUST be invoked from `sync_issues_docs.sh` so newly added items receive ATM-NNN identifiers before summary regeneration.
- §11.4.55 (Reopens-history per-item docs) — the per-item `Reopens.md` lives at `docs/issues/<ATM-NNN>/Reopens.md`.
- §11.4.57 (README.md doc-link section) — every Issues / Fixed item may be linked by ATM-NNN in external references.
- §11.4.44 (revision header) — applies to the state file's audit log if one is maintained.

Pre-build gates:

- CM-ATM-TICKET-IDS-COMPLETE — every workable-item heading in `Issues.md` AND `Fixed.md` MUST carry an `[ATM-NNN]` token matching `\[ATM-[0-9]{3,}\]`. Headings without one FAIL the gate.
- CM-ATM-TICKET-IDS-UNIQUE — no two items share an ATM-NNN. The state file's record count MUST equal the union of `Issues.md` + `Fixed.md` `[ATM-NNN]` occurrences (no orphan IDs, no duplicates).
- CM-ATM-TICKET-IDS-MONOTONIC — the maximum ATM-NNN in the state file MUST equal `record_count` (no gaps).
- CM-COVENANT-114-54-PROPAGATION — anchor literal §11.4.54 present across canonical files (`constitution/Constitution.md`, `constitution/CLAUDE.md`, `constitution/AGENTS.md`, parent consumer's `CLAUDE.md` + `AGENTS.md`).

Paired mutations (per §1.1):

- Remove `[ATM-NNN]` from one `Issues.md` heading → CM-ATM-TICKET-IDS-COMPLETE FAILs.
- Duplicate an ATM-NNN in the state file → CM-ATM-TICKET-IDS-UNIQUE FAILs.
- Delete `ATM-NNN=2` from a 5-record state file → CM-ATM-TICKET-IDS-MONOTONIC FAILs (gap at 2).
- Strip §11.4.54 literal from `constitution/CLAUDE.md` → CM-COVENANT-114-54-PROPAGATION FAILs.

No escape hatch. No `--skip-atm-assignment`, `--renumber`, or `--no-atm-id-required` flag exists. The discipline exists because unstable item identity is the exact failure mode that produces “the test was passing last week, what changed?” sessions where the operator cannot tell whether a given item is the same one tracked before or a coincidentally similarly-named successor.

Classification: universal (per §11.4.17). Applies to every project consuming the constitution submodule that maintains an Issues / Fixed split (cf. §11.4.53 scoping). Projects that do not maintain such a split are `NOT_APPLICABLE`.

§11.4.55 — Reopens-history tracking + per-item Reopens.md doc (User mandate, 2026-05-19)

Forensic anchor — verbatim user mandate (2026-05-19):

“Add a Reopens-count column to Issues / Issues_Summary / Fixed / Fixed_Summary. For any item whose reopens-count > 0, create docs/issues/ATM-NNN/Reopens.md (+ HTML + PDF) with comprehensive reopen history + Fixed cycles. Each reopen MUST include By (AI / User), On (date), Reason, Evidence, and each Fixed-marking with its reasoning chain.”

Why this anchor exists. §11.4.34 mandated that Reopened status carry a `**Reopened-Details:**` block — but that block lives in the **current** Issues.md heading and is overwritten on the next state change. A complete reopen / fix-cycle audit trail across the lifetime of an item was not previously canonical. §11.4.55 makes the full history first-class: every item that has been reopened at least once gets a dedicated Reopens.md companion doc capturing the entire chronological timeline — every reopen, every closure, every reasoning chain — so an operator / agent investigating “why was this reopened five times” has a single, authoritative source to read.

Operative rule. Every workable item with `reopens_count > 0` MUST have a companion document at `docs/issues/ATM-NNN/Reopens.md` (plus its HTML + PDF exports per §11.4.44 + §11.4.53 sync discipline). The document MUST contain:

1. **Revision header** per §11.4.44 (Revision integer + Last modified ISO 8601 UTC).
2. **Item identification** — ATM ID, Title, Type, Current Status, Current Location (Issues.md / Fixed.md), link back to the live heading.
3. **Cycle counters** — Total reopens, Total fixed cycles (items can reopen multiple times, each followed by a re-closure).
4. **Chronological timeline** — one entry per state-change event, each entry capturing:
 - **By:** AI or User (per §11.4.34 source attribution).
 - **On:** ISO date.
 - **Event:** Opened | Reopened | Fixed | Implemented | Completed.
 - **Reason:** the closed-vocabulary value from §11.4.34 (`test-failed` / `manual-testing-detected` / `captured-evidence-contradicts` / `end-user-report` / `cycle-re-discovered` / `design-reconsidered`) for reopens; captured-evidence summary for closures.
 - **Evidence:** path to or short description of captured artefact (test log path, recording path, sink-probe report, operator quote).

- **Outcome:** what the next state transition was.
- 5. **Reasoning chain for each closure** — for every Fixed / Implemented / Completed event, the reasoning chain that justified marking it closed (root-cause analysis link, captured-evidence under same-conditions per §11.4.7, gate / mutation pair that defends against regression).
- 6. **Most-recent state-change pointer** — explicit cross-reference to the current Issues.md / Fixed.md heading.

Summary-table column. Issues_Summary.md and Fixed_Summary.md MUST carry a Reopens column showing the integer count. When the count > 0 the cell MUST be a hyperlink to docs/issues/<ATM-NNN>/Reopens.md. Generators emit the column; the colorizer (§11.4.23) MAY apply a visual cue when reopens > 2 (repeated-reopen signal — investigate before closing again).

Composition:

- §11.4.34 (Reopened-Details on current heading) is the **per-event** capture; §11.4.55 is the **per-item history** aggregation. Both layers MUST be in sync — the state-machine helper updates Reopens.md on every Reopened-Details parse.
- §11.4.54 (ATM-NNN) provides the stable identifier so the per- item doc has a permanent path even after Issues→Fixed migration.
- §11.4.44 (revision header) applies to every Reopens.md.
- §11.4.45 (Status.md per-integration) — the per-item Reopens.md is the item-scoped analogue of the integration-scoped Status.md.
- §11.4.53 (Fixed_Summary parity) — Reopens column appears on BOTH summary tables symmetrically.

Pre-build gates:

- CM-REOPENS-DOC-EXISTS-WHEN-COUNT-GT-ZERO — for every state-file record with reopens_count > 0, docs/issues/<ATM-NNN>/ Reopens.md MUST exist (+ HTML + PDF). Missing doc FAILs.
- CM-REOPENS-DOC-REVISION-HEADER — every Reopens.md carries the §11.4.44 revision header.
- CM-REOPENS-COL-IN-SUMMARIES — both Issues_Summary.md and Fixed_Summary.md carry a Reopens column.
- CM-COVENANT-114-55-PROPAGATION — anchor literal §11.4.55 across canonical files.

Paired mutations (per §1.1):

- Delete docs/issues/ATM-NNN/Reopens.md for an item with reopens_count=2 → CM-REOPENS-DOC-EXISTS-WHEN-COUNT-GT-ZERO FAILs.
- Strip revision header from a Reopens.md → CM-REOPENS-DOC-REVISION-HEADER FAILs.
- Remove Reopens column from Issues_Summary.md table header → CM-REOPENS-COL-IN-SUMMARIES FAILs.
- Strip §11.4.55 literal from constitution/CLAUDE.md → CM-COVENANT-114-55-PROPAGATION FAILs.

No escape hatch. No --skip-reopens-doc, --collapse-history, --reopens-not-applicable flag. The discipline exists because loss of reopen history is the exact failure mode that produces “why does this keep happening” sessions with no auditable answer.

Classification: universal (per §11.4.17). Applies wherever §11.4.34 applies (every project tracking Reopened status).

§11.4.56 — Status_Summary parity + two-audience format (User mandate, 2026-05-19)

Forensic anchor — verbatim user mandate (2026-05-19):

“Every Status.md doc gets a Status_Summary parity companion ALWAYS in sync, exported to HTML + PDF. Two-page format: page 1 = non-developer audience (team-specific: audio team for audio Status, video team for video Status, etc.), page 2 = software engineer summary. Auto-generated after every main Status update.”

Why this anchor exists. §11.4.45 established Status.md per integration domain — but Status.md targets engineers (cites gate names, §-letter references, commit hashes, captured-evidence file paths). Non-developer stakeholders (audio team, video team, QA leads, product reviewers) need a plain-language version of the same content, **always derived** from the engineering Status.md so the two views can never diverge. §11.4.56 introduces the symmetric companion doc with a two-audience layout.

Operative rule. For every docs/<domain>/<integration>/ Status.md that satisfies §11.4.45, a companion docs/<domain>/<integration>/ Status_Summary.md MUST exist with:

1. **Revision header** per §11.4.44.
2. **Page 1 — For the <team>** (audience-specific heading; the team name is derived from the integration domain — audio team for docs/dolby/*, video team for docs/video/*, etc.). Page 1 contains:
 - Plain-language summary of current state (1-3 paragraphs).
 - **What works** (1-3 bullets, end-user-visible behaviour).
 - **What's broken or pending** (1-3 bullets).
 - **Operator / team actions** if any.
 - NO code references, NO §-letter jargon, NO captured-evidence file paths, NO gate / mutation names.
3. **Page 2 — For software engineers** — same content density as Status.md but condensed:
 - §-letter references, gate names, commit hashes, captured- evidence file paths.
 - Cross-references to Issues.md / Fixed.md items by ATM-NNN (§11.4.54).
4. **HTML + PDF exports** travel with the markdown (identical mtimes within sync wrapper granularity).
5. **Auto-generation** — the canonical helper scripts/testing/ generate_status_summary.sh <Status.md path> produces both pages. The generator MUST be invoked from a sync wrapper (e.g. scripts/testing/ sync_integration_status.sh per §11.4.45) every time the source Status.md is updated.

Page-1 audience derivation. The generator maps domain → audience name via a project-controlled mapping table (e.g. audio → Audio team, video → Video team, network → Network team, drm → DRM / streaming team). Unmapped domains default to Project team. The mapping table lives in the project's generator config, NOT in the Constitution (project-specific data per §11.4.17).

Composition:

- §11.4.45 (Status.md per integration) — Status_Summary.md COMPLEMENTS, never replaces. Both files exist together.
- §11.4.12 + §11.4.53 (parity discipline) — Status_Summary follows the same regeneration / HTML+PDF / mtime-alignment pattern.
- §11.4.44 (revision header) — applies to Status_Summary.md.
- §11.4.23 (colorizer) — Status_Summary.html receives the same type/status visual cues if it embeds tracked-item references.
- §12.10 (CONTINUATION.md) — non-developer stakeholders read Status_Summary; engineers read Status.md + CONTINUATION.md + Issues.md.

Pre-build gates:

- CM-STATUS-SUMMARY-EXISTS-FOR-EVERY-STATUS — for every docs/**/ Status.md, the sibling Status_Summary.md MUST exist (+ HTML + PDF) with $\text{mtime} \geq \text{Status.md mtime} - \text{tolerance}$.
- CM-STATUS-SUMMARY-TWO-AUDIENCE — every Status_Summary.md contains BOTH the Page 1 – For the <team> heading AND the Page 2 – For software engineers heading. Missing either FAILs.
- CM-STATUS-SUMMARY-REVISION-HEADER — §11.4.44 header present.
- CM-COVENANT-114-56-PROPAGATION — anchor literal across files.

Paired mutations (per §1.1):

- Delete a Status_Summary.md → CM-STATUS-SUMMARY-EXISTS-FOR-EVERY-STATUS FAILs.
- Remove Page 1 heading from a Status_Summary.md → CM-STATUS-SUMMARY-TWO-AUDIENCE FAILs.
- Strip §11.4.56 literal from constitution/CLAUDE.md → CM-COVENANT-114-56-PROPAGATION FAILs.

No escape hatch. No --skip-status-summary, --engineer-only, --no-audience-split flag. Plain-language stakeholder access is not negotiable — the project ships to humans, not just to engineers.

Classification: universal (per §11.4.17). Applies wherever §11.4.45 applies (every project that maintains domain-scoped Status.md docs).

§11.4.57 — README.md doc-link section + revision metadata (User mandate, 2026-05-19)

Forensic anchor — verbatim user mandate (2026-05-19):

“Add a doc-link section to README.md — links to Issues + Issues_Summary + Fixed + Fixed_Summary + CONTINUATION + ALL Status docs + their exports. Each link shows revision + last-modified.”

Why this anchor exists. Operators, AI agents, and external reviewers arriving at the project repo for the first time need a single discoverable entry point listing every canonical tracked- items + status document, along with each document’s freshness (revision + last-modified). Without this surface, agents resort to `find docs -name '*.md'` which returns hundreds of files unranked. §11.4.57 makes the README.md the canonical index.

Operative rule. Every project’s top-level README.md MUST contain a section titled Tracked-Items + Status Documents (or the project’s localized equivalent — the section heading MUST contain the literal Tracked-Items for gate detection). The section MUST be a markdown table with columns: Document (human- readable name), Last modified (ISO 8601 UTC from the doc’s §11.4.44 revision header), Revision (integer from same header), Markdown (relative link to .md), HTML (relative link to .html), PDF (relative link to .pdf).

The section MUST link to:

1. Issues.md, Issues_Summary.md (§11.4.12, §11.4.15, §11.4.16).
2. Fixed.md, Fixed_Summary.md (§11.4.19, §11.4.53).
3. CONTINUATION.md (§12.10).
4. **Every** docs/**/Status.md and its Status_Summary.md pair (§11.4.45, §11.4.56). Status docs are auto-discovered by the generator via `find docs -name 'Status.md'`.

Generator. scripts/testing/update_readme_doc_links.sh is the canonical helper. It:

1. Scans the canonical doc paths.
2. Extracts each doc’s Revision + Last modified fields from its §11.4.44 header.
3. Renders the markdown table with current values.
4. Replaces the section between explicit markers (`<!-- doc-link-section:begin -->` and `<!-- doc-link-section:end -->`) in README.md.

The wrapper MUST be invoked from `sync_issues_docs.sh` (so every Issues.md / Fixed.md update refreshes the README freshness data) AND from `sync_integration_status.sh` (so every Status.md update likewise refreshes README).

Composition:

- §11.4.12 + §11.4.19 + §11.4.53 (Issues / Issues_Summary / Fixed / Fixed_Summary parity) — README links to all four.
- §11.4.44 (revision header) — README pulls Revision + Last modified from each doc’s header.
- §11.4.45 (Status.md) + §11.4.56 (Status_Summary.md) — README enumerates every Status pair.
- §12.10 (CONTINUATION.md) — explicitly linked from README so operators discover it on arrival.

Pre-build gates:

- CM-README-DOC-LINK-SECTION-PRESENT — README.md contains the literal Tracked-Items heading AND both section markers (<!-- doc-link-section:begin -->, <!-- doc-link-section:end -->).
- CM-README-DOC-LINK-ROWS-COMPLETE — every canonical doc (Issues, Issues_Summary, Fixed, Fixed_Summary, CONTINUATION, every Status.md + Status_Summary.md found under docs/) appears as a row inside the section.
- CM-README-DOC-LINK-FRESHNESS — the Last modified value in each row matches the corresponding doc's §11.4.44 header within sync-wrapper granularity (no row staler than the source doc by more than the sync wrapper's tolerance).
- CM-COVENANT-114-57-PROPAGATION — anchor literal across files.

Paired mutations (per §1.1):

- Strip the section markers from README.md → CM-README-DOC-LINK-SECTION-PRESENT FAILs.
- Remove the CONTINUATION.md row from the section → CM-README-DOC-LINK-ROWS-COMPLETE FAILs.
- Backdate a Last modified value in a row → CM-README-DOC-LINK-FRESHNESS FAILs.
- Strip §11.4.57 literal from constitution/CLAUDE.md → CM-COVENANT-114-57-PROPAGATION FAILs.

No escape hatch. No --skip-readme-doc-links, --collapse-status-rows, --no-freshness-check flag. The discipline exists because invisible docs are non-existent docs from a discoverability standpoint, and the §11.4 covenant specifically prohibits surfaces where claims (here: doc freshness) can drift unnoticed.

Classification: universal (per §11.4.17). Applies to every project that maintains an Issues / Fixed split or Status.md docs (i.e., to every project consuming this Constitution that has any of §11.4.15 / §11.4.45 / §11.4.53 in scope).

§11.4.58 — Parallel-development methodology (User mandate, 2026-05-19)

Forensic anchor — verbatim user mandate (2026-05-19T~05:00Z MSK):

“We MUST DO one comprehensive research and planning in the background in parallel with current mainstream work: our current methodology of the development is very slow. It takes us days to fix a bunch of bugs, add some changes or features and all this to be tested, verified and shipped. From iteration to iteration we are slower and slower. We MUST CREATE adjusted improved version of working methodology where multiple workable items could be done in parallel, all come into the central (main) branch as they are done, then we at particular moment rebuild and reflash the System and in background full testing with validation and verification is done. Each parallel work on one workable (or more) item(s) must use parallel agents as much as possible! We MUST HAVE proper synchronization mechanism between parallel working routines so they do not cause conflicts, break features or create any other types of the problems! Document everything - the whole plan, whole methodology, working flow with in depth diagrams and graphs,

all user guides and manuals and then add this all into our root (constitution Submodule) Constitution, CLAUDE.MD and AGENTS.MD! Make sure we start using this ASAP since we MUST increase efficiency and productivity a couple of times now! Writing in depth tests (all supported types of the tests) with Challenges and full HelixQA use is MANDATORY! Every test we execute besides executed with success MUST RESULT in proof that actual functionality being tested REALLY DOES WORK with NO BLUFF of any kind! Heavt enforcement of no-bluff / anti-bluff policy IS MANDATORY!”

Why this anchor exists. Sequential phase-chain working pattern (Phase 39.A → Phase 39.B → ... → Phase 39.DT, observed 2026-05-05 to 2026-05-19: 86 unique Phase 39.X sub-phases across 14 days = ~6/day sub-phase rotation, 224 commits / 14 days = ~16/day) serialises on five distinct bottlenecks: (B1) `commit_all.sh` flock at `.git/.commit_all.lock`, (B2) single-thread sub-phase chaining via `docs/CONTINUATION.md` hand-offs, (B3) rebuild-blocking on every `REQUIRES_REBUILD` fix (§12.7 -j2 cap = 4–6 h rebuild + 30 min flash + 60 min validate per cycle), (B4) single-thread `test_all_fixes.sh` sweep (~30–60 min × 8 phase-device combinations), (B5) subagent dispatch latency (one-at-a-time spawn + wait pattern). Net effect: single-item wall-clock 1–4 days where productive work is < 4 h. The User mandate requires the methodology to multiply throughput several-x.

Operative rule. Project work proceeds through the Parallel Work Unit (PWU) pipeline rather than sequential phase-chain. Each PWU is a self-contained workable item with mandatory components (ATM-NNN identifier per §11.4.54, Issues.md entry per §11.4.15+§11.4.16, file- scope manifest, §11.4.43 RED test, source patch, pre-build gate per §11.4.4(b) layer 1, post-flash test per §11.4.4(b) layer 3, paired §1.1 meta-test mutation, §11.4.4(b) layer 4 HelixQA Challenge bank entry, captured-evidence directory per §11.4.5+§11.4.52). PWUs execute through five pipeline stages: **Stage 1 DEVELOP** (parallel, N PWU agents in isolated worktrees), **Stage 2 MERGE** (serial, conductor processes one PWU at a time via `commit_all.sh` flock + §11.4.41 4- step merge-first pipeline), **Stage 3 REBUILD+FLASH** (parallel where possible: AOSP build inside §12.7 bounded scope, D3+D4 flash serialised on USB hardware), **Stage 4 VALIDATE** (parallel on D3+D4+meta-test+ coverage-audit), **Stage 5 SWEEP** (parallel: HelixQA Challenge bank, Issues→Fixed migration, README.md doc-link refresh per §11.4.57, HTML+PDF exports). Stage 1 of round N+1 starts WHILE Stages 4+5 of round N are still running — this is the primary throughput multiplier.

Synchronization mechanism. Four-layer lock hierarchy: **L1** parent serial flock `.git/.commit_all.lock` (held only during Stage 2, released between PWUs), **L2** per-submodule git operations (implicit via cascade), **L3** contention-path advisory locks at `qa-results/pwu-locks/<sha256(path)>.lock` for the 10 forbidden cross- PWU paths (`CLAUDE.md/AGENTS.md/Constitution.md/CONTINUATION.md/ Issues.md/Fixed.md/pre_build_verification.sh/meta_test_false_positive_proof.sh/build.sh/commit_all.sh+push_all.sh/device.mk+BoardConfig.mk/ atmosphere-.sh/init.rc`), **L4** per-PWU git worktree (zero cross- PWU file conflicts during Stage 1). Disjoint-scope PWUs (different test files, different APKs, different drivers, different framework services) run fully parallel. Conflict detection at Stage 2 entry: `git fetch --all --prune --tags + git diff main...HEAD -- <scope> + cross-PWU scope overlap check`. Overlap detected → PWU rejected to Stage 1 for rebase (per §11.4.41 step 2 integration-before-force-push).

Anti-bluff enforcement. Conductor REFUSES to merge a PWU lacking all four anti-bluff checks: **C1** §11.4.43 RED test captured-evidence (file exists + non-zero exit code in log + timestamp predates source patch — proves test catches regression), **C2** §1.1 paired meta-test mutation (applied + gate FAILs under mutation + restored + gate PASSes — proves gate is not a bluff gate), **C3** §11.4.50 deterministic-consistency (3 iterations for normal, 10 for cycle-validation, ALL identical exit codes AND identical evidence-hashes — eliminates §11.4.7 flake/transient/intermittent demotion path), **C4** §11.4.5 captured-evidence (audio: WAV with non-trivial RMS + ffprobe channel count; video: screen recording with ffprobe frame count > 0 + analyzer matched event; UI: uiautomator dump with under-test element state; or HelixQA result.json PASS verdict + positive-evidence chain). HelixQA coverage is MANDATORY for every user-visible PWU per User mandate — Challenge bank entry in tools/helixqa/banks/atmosphere.yaml referencing the PWU's ATM-NNN + dispatching to the on-device test + scoring PASS only on positive captured evidence. Metadata-only PASS / configuration-only PASS / absence-of-error PASS / grep-without-runtime-evidence PASS all REJECTED at the merge gate.

Agent orchestration. Four roles: **Conductor** (1, main thread) — issues ATM-NNN, dispatches PWU agents, runs Stage 2 merge serially, runs Stages 3-5, owns all contention-path edits, maintains docs/CONTINUATION.md §3 live snapshot. **PWU agent** (4-6 background) — one PWU each, Stage 1 work in isolated worktree, signals qa-results/pwu-<atm>/READY_FOR_MERGE on completion. **Validator** (4 background) — Stage 4 work in parallel on D3+D4+meta-test+coverage-audit. **Sweeper** (2 background) — Stage 5 work in parallel with next round's Stage 1. Max concurrent agent count ~12 at peak, bounded by §12.6 60% memory budget (~6 GB total ≤ 38 GB budget). Composition with §12 host-session safety: conductor refuses dispatch if host_check_safety fails; per §12.7 AOSP build still hard-capped at -j2 inside bounded scope; per §12.8 + §12.8b + §12.8c concurrency gates still block on stray gradle/kotlin daemons.

Composition map (mandatory):

- §11.4.4 — four-layer test coverage per PWU (pre-build + post-build + on-device + HelixQA bank).
- §11.4.5 — audio/video quality analysis comprehensiveness (PWU captured-evidence directory).
- §11.4.6 — no-guessing mandate (every PWU claim backed by captured forensic evidence or explicit UNCONFIRMED: tag).
- §11.4.7 — demotion-evidence rule (3/3 deterministic-consistency closes the demotion-bluff path).
- §11.4.9 — batch-source-fixes-before-rebuild (Stage 2 → Stage 3 batch threshold).
- §11.4.15 + §11.4.16 + §11.4.33 — PWU status + type vocabulary + type-aware closure terminal value.
- §11.4.19 — atomic Issues.md → Fixed.md migration on PWU closure.
- §11.4.41 — Pre-Force-Push Merge-First 4-step pipeline at Stage 2 entry.
- §11.4.42 — Iteration-discipline (PWU pipeline IS the iteration conductor; §11.4.58 binds steps 1-5 to PWU stages 1-5).
- §11.4.43 — TDD RED-first per-PWU enforced at merge time.
- §11.4.45 — Integration-status-doc maintenance (PWU's Status.md updated at every stage transition).
- §11.4.49 — Dual-approach testing (PWU MUST ship both UI-driven AND Intent/Broadcast-driven variants when applicable).

- §11.4.50 — Deterministic-consistency (3-iter normal / 10-iter cycle-validation merge-time enforcement).
- §11.4.52 — Autonomous-validation (PWU's on-device test MUST run without operator presence per Stage 4 parallel design).
- §11.4.54 — ATM-NNN ticket identifier (PWU identifier replaces Phase 39.X letters).
- §11.4.57 — README.md doc-link refresh at Stage 5.
- §12.6 — 60% memory-budget ceiling (caps concurrent agent count).
- §12.7 — AOSP m -j hard-cap at -j2 (Stage 3 build cap).
- §12.8 + §12.8b + §12.8c — concurrency hardening (Stage 3 preflight refuses on running gradle/kotlin/git-pack-objects/load-per-cpu spike).
- §12.10 — CONTINUATION.md maintenance (conductor owns the file by §3.4 forbidden cross-PWU rule).
- §9.2 — data safety (every history rewrite during Stage 2 merge-first step 2 backed by hardlinked .git backup).

Pre-build gates:

- CM-PWU-LOCK-HIERARCHY — verify lock hierarchy script exists at scripts/testing/pwu_acquire_path_lock.sh and scripts/testing/pwu_release_path_lock.sh, both executable, both reference the 10 forbidden cross-PWU paths from §3.4 of docs/guides/PARALLEL_DEVELOPMENT_METHODODOLOGY.md.
- CM-PWU-ANTI-BLUFF-COVERAGE — for every PWU with status “Ready for merge” or beyond, assert all four anti-bluff checks (C1–C4) have captured-evidence files under qa-results/pwu-<atm>/evidence/. PWU lacking any of red-test-output / mutation_patch / deterministic-consistency / captured-feature-proof FAILs the gate.
- CM-PWU-MERGE-QUEUE-DISCIPLINE — scan qa-results/pwu-queue/ for PWUs that bypassed the merge queue (commits on main without a corresponding qa-results/pwu-queue/merged/<atm> marker). Any bypass FAILs the gate.
- CM-PWU-PARALLEL-AGENT-LIMIT — assert no more than 6 concurrent background PWU agents (count tmux/screen sessions or worktree branches) — over-limit FAILs the gate (per §12.6 memory budget).
- CM-COVENANT-114-58-PROPAGATION — anchor literal §11.4.58 present in constitution/CLAUDE.md + constitution/AGENTS.md + parent CLAUDE.md/AGENTS.md + every owned-submodule CLAUDE.md/AGENTS.md (42 files total in current ATMOSphere consumer family).

Paired mutations (per §1.1):

- Move scripts/testing/pwu_acquire_path_lock.sh aside → CM-PWU-LOCK-HIERARCHY FAILs.
- Touch a PWU's READY_FOR_MERGE marker without red-test-output.log → CM-PWU-ANTI-BLUFF-COVERAGE FAILs.
- Land a commit on main without a corresponding qa-results/pwu-queue/merged/<atm> marker → CM-PWU-MERGE-QUEUE-DISCIPLINE FAILs.
- Spawn a 7th concurrent PWU agent → CM-PWU-PARALLEL-AGENT-LIMIT FAILs.
- Strip §11.4.58 literal from constitution/CLAUDE.md → CM-COVENANT-114-58-PROPAGATION FAILs.

No escape hatch. No `--skip-merge-queue`, `--allow-bypass`, `--no-anti-bluff-check`, `--unlimited-agents`, `--sequential-phase-chain-mode` flag exists in any pipeline helper. The discipline exists because (a) the User mandate explicitly requires the multiplier ASAP and (b) the §11.4 covenant specifically prohibits PASS-bluffs (which the four anti-bluff merge-time checks mechanically prevent). Operators who feel they need to bypass the pipeline for a specific high-urgency fix should split the work into a single trivial PWU and let the pipeline run normally — Stage 1 of a trivial PWU is < 30 min, Stage 2 is < 5 min, Stage 3 is the same rebuild cost they would pay anyway.

Classification: universal (per §11.4.17). Applies to every project consuming this Constitution that has multiple owned-submodules + a sequential commit/push tool + a captured-evidence testing discipline. Project-specific implementations (agent dispatch helpers, worktree layout, batch thresholds) live in consumer-side docs/guides/`PARALLEL_DEVELOPMENT_METHODODOLOGY.md` (this Constitution defines the discipline; consumers implement the tooling).

Reference: parent `docs/guides/PARALLEL_DEVELOPMENT_METHODODOLOGY.md` (full methodology with diagrams, templates, operator manual, and migration plan).

§11.4.59 — README always-sync mandate (User mandate, 2026-05-19)

Forensic anchor — direct user mandate (verbatim, 2026-05-19):

“fully review and update our main README document. Some points are not valid anymore, some are missing. Make sure main README is among documents we MUST ALWAYS keep updated and in Sync with the projects and other documentation! Make sure we always export it (on every update) into PDF and HTML. This mandatory rules/constraints MUST BE all added into the root (constitution Submodule) Constitution, AGENT.MD and CLAUDE.MD!”

`README.md` at the project root is a §11.4.12-class always-sync document. It MUST be:

1. **Reviewed and updated whenever** new docs / integrations / Status.md entries appear, new submodules land, applied-fixes count changes, or canonical paths shift.
2. **Kept in lockstep** with docs/CONTINUATION.md (§12.10) and the Issues / Issues_Summary / Fixed / Fixed_Summary doc set (§11.4.12, §11.4.53) — same always-sync discipline.
3. **Exported to .html and .pdf on every update.** New helper scripts/testing/sync_readme_export.sh performs the pandoc + weasyprint export; it is auto-invoked by scripts/testing/sync_issues_docs.sh so a single doc-sync run refreshes the entire doc surface (Issues / Issues_Summary / Fixed / Fixed_Summary / CONTINUATION / README — md + html + pdf).
4. **Carry a §11.4.44 revision header** (Revision N + Last modified ISO timestamp) at the top of the file.
5. **Contain a Documentation Map section** linking to every Status.md
 - Status_Summary.md + spec + plan + guide + script-companion doc + changelog + the constitution submodule, plus a per-audience navigation (end user / developer / QA / agent).

6. **Self-contained** — no hyperlinks to ephemeral external systems as the only source of truth.

Stale exports are §11.4.59 violations regardless of whether the underlying README.md is correct — an operator (or future agent) reading the HTML or PDF gets a divergent view, and the §12.10 CONTINUATION resumption guarantee silently breaks. Same discipline as §11.4.12 Issues_Summary and §11.4.53 Fixed_Summary applied to README.md.

Captured-evidence enforcement. Pre-build gate CM-README-EXPORT-SYNC locks four invariants: (a) README.md exists, (b) README.html exists, (c) README.html mtime ≥ README.md mtime, (d) README.pdf mtime ≥ README.md mtime (skipped gracefully if weasyprint is unavailable on the host but enforced when the PDF is present). Paired meta-test mutation backdates README.html + README.pdf so the source is newer → gate FAILs.

Composition. Composes with §11.4.12 (Issues_Summary parity), §11.4.18 (script-companion docs), §11.4.23 (visual-cue HTML colorizer not applicable to README's narrative format), §11.4.44 (revision header on every Status.md and now README.md), §11.4.45 (Status.md integration), §11.4.53 (Fixed_Summary parity — README is the canonical sibling at the project-overview layer), §11.4.56 (Status_Summary parity), §11.4.57 (README.md is the canonical index per Phase 39.EW+1), §12.10 (CONTINUATION.md resumption guarantee depends on README's Documentation Map being current).

No escape hatch. §11.4.59 has NO operator-facing override flag. No --skip-readme-sync, --no-readme-export, --readme-stale-OK flag exists. The discipline exists because README.md is the first file any new agent / operator / contributor reads — letting it diverge from reality is the exact §11.4 PASS-bluff pattern but at the project-overview layer.

Classification: universal (per §11.4.17). Applies to every project consuming this Constitution. The doc-sync helper, pandoc + weasyprint dependencies, and gate name are universal; the specific Documentation Map contents are consumer-side per-project.

Canonical authority: constitution submodule Constitution.md §11.4.59.

Non-compliance is a release blocker regardless of context.

§11.4.60 — Documentation always-sync composite covenant (User mandate, 2026-05-19)

Forensic anchor — verbatim user mandate (2026-05-19 ~09:00Z):

“Double check if all documents are properly tied with our root Constitution, CLAUDE.MD and AGENTS.MD so they are always up to date, always in sync and exported into PDF and HTML! ... Issues, Issues_Summary, Fixed, Fixed_Summary, Continuation, Status and Status_Summary for all contexts (areas) — THEY ALL MUST BE REGULARLY UPDATED, IN SYNC AND CONSISTENT without giving at any moment false picture about the state of the project or particular area(s) of it!”

The covenant. Eight documentation classes constitute the project’s living state surface. They MUST be in sync at all times across markdown + HTML + PDF artefacts. Per-class anchors §11.4.12 / §11.4.44 / §11.4.45 / §11.4.53 / §11.4.56 / §11.4.57 / §11.4.59 / §12.10 each govern one class; §11.4.60 binds them via a single mechanically-enforced composite invariant so the failure mode “one per-class gate silently disabled while operator reads a divergent HTML/PDF” is structurally impossible.

Eight bound doc classes (artefact triple — .md + .html + .pdf):

1. docs/Issues.md — open-item tracker (§11.4.12 + §11.4.15 + §11.4.16)
2. docs/Issues_Summary.md — auto-generated short form (§11.4.12)
3. docs/Fixed.md — closed-item archive (§11.4.53)
4. docs/Fixed_Summary.md — auto-generated short form (§11.4.53)
5. docs/CONTINUATION.md — resumption handoff (§12.10)
6. README.md — project overview + doc-link index (§11.4.57 + §11.4.59)
7. docs/**/Status.md (every domain-scoped instance) — per-integration status (§11.4.45)
8. docs/**/Status_Summary.md (every domain-scoped instance) — auto-generated short form (§11.4.56)

Mandatory composite gate. CM-DOCS-COMPOSITE-SYNC (pre-build, device/rockchip/rk3588/tests/pre_build_verification.sh):

- For every doc-class instance with .md present, verify .html mtime $\geq$.md mtime AND .pdf mtime $\geq$.md mtime.
- For every Status.md instance, also verify the sibling Status_Summary.md mtime $\geq$ Status.md mtime (if the summary exists).
- Walks docs/ recursively for Status.md — the Status fleet is not enumerated statically because new integrations land per phase.

Composite gate FAILs the build if ANY single instance fails. The per-class gates (CM-ISSUES-SUMMARY-SYNC, CM-DOCS-EXPORT-SYNC, CM-FIXED-SUMMARY-SYNC, CM-CONTINUATION-DOC-INSYNC, CM-README-EXPORT-SYNC) remain in place — §11.4.60 is the belt-and-suspenders layer that catches what they miss.

Paired mutation (per §1.1). meta_test_false_positive_proof.sh backdates docs/Issues.html to year 2000 (touch -t 200001010000.00) and asserts CM-DOCS-COMPOSITE-SYNC FAILs.

Revision-header coupling. Per §11.4.44 every one of the 8 doc classes carries the ****Revision:** N + **Last modified:**** ISO 8601 UTC header directly under the H1. CM-DOCS-COMPOSITE-SYNC does NOT re-check the revision header (already covered by CM-DOC-REVISION-HEADER-PRESENT) — composite gate is artefact-mtime only, by design narrow.

Auto-sync wrapper coupling. Per §11.4.12 every edit to an Issues/Fixed-class doc flows through sync_issues_docs.sh. Per §11.4.45 every edit to a Status.md flows through sync_integration_status.sh <Status.md>. Per §11.4.59 every edit to README.md flows through sync_readme_export.sh. §11.4.60 does NOT change the wrapper contract — it only catches the failure mode where an operator (or AI agent) bypassed the wrapper and committed markdown without regenerating exports.

No escape hatch. No `--skip-composite-doc-sync`, `--allow-stale-html`, `--summary-not-applicable` flag exists. The covenant exists for the operator's own protection — the moment a single per-class gate is bypassed, divergence accumulates silently until release-time forensics finds it (the exact §11.4 PASS-bluff pattern the canonical covenant was authored to eliminate).

Composes with §11.4.12 (Issues_Summary), §11.4.15 (Status field), §11.4.16 (Type field), §11.4.19 (atomic Issues→Fixed migration), §11.4.23 (colorizer), §11.4.33 (type-aware closure vocabulary), §11.4.44 (revision header), §11.4.45 (Status.md), §11.4.53 (Fixed_Summary), §11.4.56 (Status_Summary), §11.4.57 (README doc-link), §11.4.59 (README always-sync), §12.10 (CONTINUATION maintenance).

Propagation. Pre-build gate `CM-COVENANT-114-60-PROPAGATION` enforces the §11.4.60 anchor literal in every `CLAUDE.md` / `AGENTS.md` across parent + 10 owned submodules + HelixQA dependencies (same shape as existing §11.4.5X propagation gates). The propagation gate is `OPTIONAL` for Phase 39.FC's initial landing — only the composite gate + its paired mutation are mandatory; propagation gates trail per established phase pattern.

Classification: universal (per §11.4.17). Applies to every project consuming this Constitution that maintains any of the 8 doc classes. The composite gate logic is universal; the specific doc-class instance paths are consumer-side per-project.

Canonical authority: constitution submodule `Constitution.md` §11.4.60.

Non-compliance is a release blocker regardless of context.

§11.4.61 — Mandatory Markdown metadata table + structured-doc ToC (User mandate, 2026-05-19)

Forensic anchor — direct user mandate (verbatim, 2026-05-19):

“For every Markdown document which contains structured content (with headings / sections and sub-sections) make sure that every time we apply change to the structure, table of contents on the top of the document is created or updated! This is **MANDATORY** for every structured MARKDOWN document. Automatically its PDF and HTML versions **MUST BE** (re)generated!”

“Introduce ... revision number, date and time of creation, date and time of last modification, other useful information we have in documents such as Issues, Issues_Summary, Fixed, Fixed_Summary, Status, Status_Summary, Continuation and similar. ... make this mandatory for **EVERY** Markdown document from now on, update root constitution Submodule with these changes and commit and push it all to all upstreams.”

Scope of §11.4.61. This clause adds two universal disciplines that sit alongside (NOT replacing) the §11.4.65 universal export mandate. §11.4.65 governs the existence and freshness of `.html` + `.pdf` siblings; §11.4.61 governs the *content* of the Markdown sources: how revision/status metadata is displayed at the top, and how the document's navigation map (Table of contents) is kept in lockstep with the heading structure. The two clauses compose — §11.4.65 guarantees readers see *something current*; §11.4.61 guarantees what they see communicates its own provenance and is internally navigable.

A. Canonical metadata table (supersedes §11.4.44 format). The §11.4.44 bold-line revision header is superseded by a **Markdown table** placed immediately after the document’s H1 title (and any blank line that follows). Pre-existing bold-line headers remain historically valid but **MUST** migrate to the canonical table at the next substantive edit. The canonical table has these **MANDATORY** rows:

Field	Value
Revision	positive integer; starts at 1, monotonic, never reset
Created	ISO 8601 date — first commit touching the file
Last modified	ISO 8601 date — most recent commit touching the file
Status	one of <code>active</code> / <code>draft</code> / <code>deprecated</code> / <code>superseded</code>

ENCOURAGED rows (**REQUIRED** when the document tracks workable items or has a known continuation; render with `–` / `none` when N/A):

Field	Value
Status summary	one-line current-state hook
Issues	comma-separated workable-item IDs
Issues summary	one-line summary of open issues
Fixed	comma-separated workable-item IDs of resolved items
Fixed summary	one-line summary of fixes
Continuation	next action or linked CONTINUATION.md anchor

Additional rows **MAY** be added for project-specific metadata as long as the **MANDATORY** rows are present.

B. Structured-document Table of contents. Any tracked `*.md` with **two or more H2 sections** (“structured content”) **MUST** include a `## Table of contents` section immediately after the metadata table. The ToC **MUST** list every H2 and H3 in document order with anchor links and **MUST** be regenerated on every structural change (heading added / removed / renamed / reordered). A stale ToC is a §11.4 **PASS-bluff**: operators reading the rendered Markdown see a navigation map that no longer matches the body.

Scope (IN-scope under §11.4.61). Every tracked `*.md` in the §11.4.65 **INCLUDED** set (project-root `*.md`, `docs/**/*.*.md`, `scripts/**/*.*.md` companion docs, owned-submodule trees, `constitution/**/*.*.md`, owned HelixQA dependencies). The §11.4.44 exclusions for `CLAUDE.md` / `AGENTS.md` / `README.md` are **SUPERSEDED** here too — explicit metadata in the file itself is preferred over indirect tracking via a separate `VERSION` file because operators reading the rendered document see the freshness directly.

Narrow exemptions (NOT IN-scope). Same **EXCLUDED** set as §11.4.65 (external/`**`, prebuilts/`**`, out/`**`, build/`**`, application- internal `*.md` shipped with source code, non-owned third-party submodule trees). Additional exemptions for the structural rules: `LICENSE`, `LICENSE.md`, `NOTICE`, `VERSION`, `OWNERS`, machine- generated `CHANGELOG.md`.

PDF + HTML siblings. Deferred to §11.4.65. The metadata table and ToC are *content* requirements; their export-freshness guarantee lives in §11.4.65's CM-UNIVERSAL-MARKDOWN-EXPORT-SYNC gate. A §11.4.61 metadata or ToC change is a .md modification, which transitively triggers §11.4.65 regeneration — the two clauses compose without duplicate enforcement.

Anti-bluff captured-evidence gates (planned):

- CM-MD-METADATA-PRESENT — walks every tracked *.md (minus exemptions) and asserts (a) H1 present, (b) within 25 lines below H1 a Markdown table contains all four MANDATORY rows (| Revision |, | Created |, | Last modified |, | Status |).
- CM-MD-TOC-PARITY — for every *.md with ≥ 2 H2 sections, asserts a ## Table of contents is present and its entries' anchor slugs match the document's live H2/H3 set in order.

Paired §1.1 mutation tests (planned):

- Strip the | Revision | row from one *.md → CM-MD-METADATA-PRESENT FAILs.
- Rename one H2 without updating ToC → CM-MD-TOC-PARITY FAILs.

Composition with §11.4.44. §11.4.44 remains the authoritative clause for the *fact* that documents must carry revision metadata; §11.4.61 changes the *format* (bold-lines → table) and *scope* (docs/**/*.*.md → every tracked *.md per §11.4.65 INCLUDED set). Consumer projects with existing gates that grep for ****Revision:**** MUST update those gates to also accept the canonical table form. Migration period: 30 days from §11.4.61 landing.

Composition with §11.4.45 / §11.4.59 / §11.4.60 / §11.4.65. The per-class always-sync clauses (Status.md §11.4.45, README §11.4.59, composite §11.4.60) and the universal export covenant §11.4.65 govern *artifact existence and freshness*; §11.4.61 governs *content discipline inside the source .md*. The clauses are orthogonal layers of the same anti-bluff posture.

No escape hatch. No --skip-md-metadata, --no-metadata-table, --toc-stale-OK, --allow-missing-revision flag exists. Divergent metadata and stale ToCs are §11.4 PASS-bluff patterns at the doc-surface layer.

Classification: universal (per §11.4.17). Applies to every project consuming this Constitution. The canonical table format and ToC mandate are universal; the per-project tracker IDs (e.g. HRD-) and specific Issues/Fixed-class doc paths are consumer-side per-project.

Canonical authority: constitution submodule Constitution.md §11.4.61.

Non-compliance is a release blocker regardless of context.

§11.4.63 — Workable-items procedure docs as single source of truth (User mandate, 2026-05-19)

Forensic anchor — verbatim user mandate (2026-05-19 ~10:30Z):

“To make workable items tracking and creation through all of the mentioned mandatory documents we MUST create proper procedure document which will be named in Constitution, AGENTS.md and CLAUDE.md as single source of truth for this procedure! Make sure it is created under: docs/procedures/issues/Creation.md which will be always exported to PDF and HTML whenever it is updated. We MUST create procedure documents for Updating.md, Resolution.md, Reopening.md and other workable items actions that we have. ... Everything done on project - new features, changes, tasks, bug fixes, investigation, ordinary tasks (writing documentation for example) MUST ALL go through the workable items flow and proper procedures!”

Every workable-item action — opening, updating, closing, reopening, migrating across Issues.md [?] Fixed.md — MUST follow the canonical procedure document at docs/procedures/issues/<Action>.md. The five procedure docs are MANDATORY references; no agent or operator may invent ad-hoc procedures. The closed-set:

Procedure doc	Covers
Creation.md	Opening a new workable item (Bug / Feature / Task)
Updating.md	Non-closure, non-reopen edits — status transitions, evidence append, type re-classification
Resolution.md	Atomic Issues.md → Fixed.md close (Bug→Fixed, Feature→Implemented, Task→Completed per §11.4.33)
Reopening.md	Atomic Fixed.md → Issues.md reopen with §11.4.34 source attribution + §11.4.55 reopens-history
Migration.md	Bidirectional atomic-move mechanics + sync_issues_docs.sh internals (consumed by Resolution + Reopening)

Each procedure doc carries the §11.4.44 revision header + HTML + PDF exports synchronised to its .md source. New helper scripts/testing/sync_procedure_docs_export.sh walks the procedure- docs directory and emits matching .html + .pdf for each .md; invoked as the final stage of scripts/testing/sync_issues_docs.sh so a single operator invocation refreshes every doc-side artifact.

Universality of scope. All work — new features, behavioural changes, tasks (refactor / doc / infra / gate / audit / cleanup), bug fixes, investigation, documentation work — MUST flow through these procedures. The “I’ll just tweak this one file” escape hatch is forbidden because it produces silent doc drift and reopens the §11.4 covenant’s PASS-bluff failure mode at the process-tracking layer.

Composes with §11.4.6 (no-guessing — every procedure step is mechanical, no likely / seems), §11.4.7 (demotion-evidence — Reopening.md is the canonical implementation), §11.4.11 (file-layout — procedure docs live in docs/procedures/issues/, qa-results forensics live in qa-results/), §11.4.12 (Issues_Summary sync — all five procedures invoke sync_issues_docs.sh whose pipeline is documented in Migration.md), §11.4.15 (Status closed-set — all five

procedures cite the same closed-set), §11.4.16 (Type closed-set), §11.4.19 (atomic Issues² Fixed move — Migration.md is the implementation), §11.4.23 (visual cue + grouping colorization — invoked by sync wrapper, no procedure-side bypass), §11.4.33 (type- aware closure vocabulary — Resolution.md enforces), §11.4.34 (Reopened-source attribution — Reopening.md enforces), §11.4.44 (revision header on every procedure doc), §11.4.53 (Fixed_Summary parity — Migration.md documents the pipeline), §11.4.54 (ATM-NNN identifier — Creation.md allocates), §11.4.55 (Reopens-history per- item doc — Reopening.md authors), §11.4.60 (documentation always- sync composite — procedure docs are an additional doc class binding to the composite covenant).

Pre-build gate CM-PROCEDURES-DOCS-PRESENT checks (1) all 5 procedure docs exist at docs/procedures/issues/{Creation,Updating,Resolution,Reopening,Migration}.md, (2) each carries the §11.4.44 revision header, (3) each has matching .html + .pdf exports with mtime ≥ the .md mtime, (4) scripts/testing/sync_procedure_docs_ export.sh exists + is executable, (5) scripts/testing/sync_issues_ docs.sh invokes sync_procedure_docs_export.sh. Paired mutation: rename one procedure doc → gate FAILs.

Propagation gate CM-COVENANT-114-63-PROPAGATION enforces this anchor literal in every CLAUDE.md / AGENTS.md across parent + 10 owned submodules + nested submodules + HelixQA dependencies. Paired mutation strips the anchor literal → gate FAILs.

No escape hatch — no --ad-hoc-procedure, --skip-procedure-doc, --procedure-not-applicable flag exists. Inventing an alternate procedure for a “small” change is the exact failure mode this anchor closes.

Classification: universal (per §11.4.17). The closed-set of five procedure docs and the sync-wrapper composition are universal; the specific test-name / fix-name examples within each procedure doc may be consumer-side per-project.

Canonical authority: constitution submodule Constitution.md §11.4.63.

Non-compliance is a release blocker regardless of context.

§11.4.65 — Universal Markdown export mandate (User mandate, 2026-05-19)

Forensic anchor — direct user mandate (verbatim, 2026-05-19):

“Any markdown document inside the project and which is not part of the applications or services source code MUST BE exported (be available) in PDF and HTML! Any already existing Markdown document that fulfills this condition and which does not have HTML or PDF at all or it is not in sync with it MUST HAVE (re)generated PDF and HTML version! Every time when Markdown document (file) is modified, its proper HTML and PDF versions MUST BE regenerated. Markdown documents MUST BE at all times in sync with PDF and HTML versions!”

The covenant generalizes the per-class always-sync discipline of §11.4.12 / §11.4.18 / §11.4.44 / §11.4.45 / §11.4.53 / §11.4.56 / §11.4.57 / §11.4.59 / §11.4.60 / §11.4.63 / §11.4.64 to **every Markdown document in the project that is not part**

of an application or service's source-code tree. Per-class anchors govern specific files (Issues, Fixed, Status, README, etc.); §11.4.65 is the catch- all: any *other* .md file that documents the project — guides, research notes, plans, hardware-ID notes, script-companion docs, changelogs, procedure docs, README files inside owned submodules — MUST also have .html + .pdf siblings, all three artefacts in sync at all times. Stale exports of *any* documentation surface are the exact “operator reads divergent HTML/PDF while .md is correct” PASS-bluff §11.4 forbids — only the per-class enforcement so far prevented it for specifically-named docs; §11.4.65 closes the long tail.

Scope (closed-set):

INCLUDED: - Project root *.md (README.md, CLAUDE.md, AGENTS.md, CONTRIBUTING.md, etc.) - docs/**/*.*md (guides, research, plans, changelogs, procedures, hardware notes) - scripts/**/*.*md (script-companion docs in documentation format, NOT shebang scripts) - Owned-submodule trees at device/rockchip/atmosphere/<submodule>/: top-level README.md / CLAUDE.md / AGENTS.md / CHANGELOG.md and any docs/**/*.*md - constitution/**/*.*md (the canonical-root submodule) - tools/helixqa/<helixqa-submodule>/top-level README.md / CLAUDE.md / AGENTS.md and any docs/**/*.*md (owned HelixQA dependencies)

EXCLUDED (NOT subject to §11.4.65 — these are application/service source code or third-party trees we do not own): - external/** (AOSP-mainline / upstream-mirror source trees) - prebuilts/** (binary prebuilts + their bundled docs) - packages/modules/** (AOSP mainline modules) - kernel-5.10/** (kernel source tree — has its own upstream docs) - out/** (build output) - build/** (AOSP build system) - Application / service source-code trees (e.g. Java/Kotlin module-internal *.md shipped with library code, README files inside source-only third-party gradle module directories) - Any third-party submodule NOT in the owned-submodule set (presenter, vlc-player, nova-player, mpv-player, gramophone-player, rhythm-player, strep-player, smarttube-player, torrservice, lampa) or HelixQA-owned set (Challenges, Containers, DocProcessor, LLMOrchestrator, LLMProvider, VisionEngine)

Mandatory protections (ALL must hold):

1. **Every INCLUDED .md file has .html and .pdf siblings.** A missing export is a §11.4.65 violation regardless of when the markdown was last touched.
2. **.html and .pdf mtime ≥ .md mtime** (within the same sync-wrapper invocation granularity). Stale exports are violations even if the .md itself is correct.
3. **Every modification triggers regeneration.** Whether via direct sync-wrapper invocation (sync_all_markdown_exports.sh), a git pre-commit hook auto-regenerating on staged .md changes, or the existing per-class wrappers (§11.4.12 sync_issues_docs.sh, §11.4.18 script-docs sync, §11.4.59 sync_readme_export.sh, etc.) all of which delegate the universal export path to the same canonical helper.
4. **Pre-build gate enforces parity.** CM-UNIVERSAL-MARKDOWN-EXPORT- SYNC walks the INCLUDED scope, verifies every .md has .html
 - .pdf siblings with mtime ≥ .md mtime, and FAILs the build if any are missing or stale.
5. **No escape hatch.** No --skip-md-exports, --no-pdf-only, --md-export-not-applicable, --application-internal-doc flag exists for files inside

the INCLUDED scope. The EXCLUDED scope is the only legitimate path to opt out, and it is closed-set.

Canonical helper. `scripts/testing/sync_all_markdown_exports.sh` walks the INCLUDED scope, invokes pandoc (HTML) + weasyprint (PDF) with `timeout 60` each (graceful per-file degradation), uses `docs/_progress-style.css` for visual consistency, and supports `--check-only` (exit nonzero if any out-of-sync, print list) and `--regenerate-all` (force) modes. Caps at 500 candidates with explicit abort+list if the scope is unexpectedly large (e.g. after a new submodule lands without scope-update review). Idempotent.

Composition. Composes with §11.4.12 (Issues_Summary sync — universal helper is the back-end), §11.4.18 (script-companion docs), §11.4.23 (HTML colorizer continues to post-process the always-sync docs for visual cues), §11.4.44 (revision header on every .md the universal helper exports), §11.4.45 (Status.md auto-sync), §11.4.53 (Fixed_Summary), §11.4.59 (README always-sync — the universal helper covers the long tail of every *other* root-level .md too), §11.4.60 (composite always-sync covenant — §11.4.65 is the structural generalization), §11.4.63 (procedure docs — already always-sync; §11.4.65 adds catch-all for every doc class not yet anchored), §11.4.64 (topic discoverability — summary docs the universal helper exports continue to be post-processed by `inject_topic_index.py`).

Pre-build gates:

- CM-UNIVERSAL-MARKDOWN-EXPORT-SYNC — invariants: (a) `scripts/testing/sync_all_markdown_exports.sh` exists + executable, 1. running it in `--check-only` mode returns 0 (no out-of-sync file in the INCLUDED scope).
- CM-COVENANT-114-65-PROPAGATION — anchor literal 11.4.65 in every CLAUDE.md / AGENTS.md across canonical-root + parent + 10 owned submodules + nested submodules + HelixQA dependencies (42-file fleet — same propagation set as §11.4.58).

Paired meta-test mutations: one strips the CM-UNIVERSAL-MARKDOWN-EXPORT-SYNC gate's enforcement literal; one strips the §11.4.65 anchor literal from a sentinel propagation file. Both mutations assert the corresponding gate FAILs when applied.

Classification: universal (per §11.4.17). The mandate to export every project-doc Markdown to HTML+PDF is universal across every project that consumes this Constitution; the specific INCLUDED / EXCLUDED scope path lists are consumer-side per-project (since each consumer has different submodule topology, build trees, and source- code roots).

Canonical authority: constitution submodule [Constitution.md](#) §11.4.65.

Non-compliance is a release blocker regardless of context.

§11.4.66 — Blocker-resolution interactive-clarification mandate (User mandate, 2026-05-19)

Forensic anchor — direct user mandate (verbatim, 2026-05-19):

“If any blockers which can be resolved with interactive response ever happen again, perform in depth research on options doable by your side and how much inputs from us you really need, then create options and present them to us. After we answer, preferably you will be unblocked and be able to continue work on blocked items. Let us make this main approach when such situations (blockers) do happen!”

When any task is blocked (waiting on operator decision, hardware access, external-service authorization, ambiguous scope, or any input the agent cannot mechanically derive), the agent **MUST** follow this five-step discipline before idling or asking free-form questions:

1. **In-depth research on agent-side options.** Identify the maximum scope that can land unilaterally without operator input — research the codebase, upstream documentation, existing tooling, the current state, what files/components are involved, what tests/gates already exist. Surface every agent-actionable path even if it requires a small operator confirmation.
2. **Calculate minimum-viable operator input.** Reduce the question set to the smallest closed set of operator-only decisions (preference, authorization, scope-bounding, hardware availability, schedule). Anything that can be researched **MUST NOT** be asked.
3. **Construct 2–4 mutually-exclusive options.** Each option carries
 1. a short label, (b) a one-line trade-off description, (c) an explicit statement of what the agent will do *after that answer*. One option marked “Recommended” with rationale. Options **MUST** genuinely differ — restating one option as three close variants is itself a §11.4.66 violation.
4. **Present via the platform’s interactive question mechanism.** On Claude Code that is `AskUserQuestion` (max 4 questions, each with 2–4 options; supports multi-select for non-mutually-exclusive sets). Other platforms (Copilot CLI, Codex, Gemini CLI) use their equivalents. **NEVER** inline free-text “what would you like?” when interactive options would do — that wastes operator attention.
5. **After the operator answers, resume work without additional round-trips.** Every option’s promised action **MUST** be sufficient to unblock — if a follow-up clarifying question is needed, the prior options were insufficiently researched and that’s itself a §11.4.66 violation. The contract is: ask once, unblock, continue.

Composes with:

- §11.4.6 (no-guessing — interactive options replace agent guessing about operator preference)
- §11.4.7 (demotion-evidence — operator preference is captured- evidence for direction changes)
- §11.4.40 (full-suite retest — uninterrupted work post-answer means the test cycle resumes cleanly without context drift)
- §11.4.41 (merge-first — when operators make conflicting choices across sessions, the merge-first discipline preserves both)
- §11.4.42 (iteration-discipline — interactive-question loop is built into the per-cycle methodology)

- §11.4.52 (autonomous-validation — most blockers can be researched to autonomous-path level, leaving only true operator preference to interactive-ask)

Mandatory protections:

- **No silent waiting.** If blocked, ask. If genuinely no operator input is needed, do not ask — proceed.
- **No bulk-text questions when interactive options would do.** Free-text prompts are reserved for truly open-ended ideation, never for closed-set decisions.
- **Each interactive question minimizes operator cost** — the operator should be able to answer all questions in under 30 seconds total.
- **Recommended option is highlighted explicitly** — the operator should know what the agent would default to if no answer comes.
- **Out-of-band channels (slack, voice, email) NEVER substitute** — the interactive mechanism creates an audit trail that free-text channels do not.

Pre-build gate CM-COVENANT-114-66-PROPAGATION enforces the anchor literal is present across the 42-file consumer fleet (parent CLAUDE.md / AGENTS.md + 10 owned-submodule pairs + HelixQA-owned pairs). Paired meta-test mutation strips the literal → gate FAILs. No escape hatch — no `--skip-ask`, `--silent-wait`, `--free-form-only` flag is permitted.

Canonical authority: constitution submodule [Constitution.md](#) §11.4.66.

Non-compliance is a release blocker regardless of context.

§11.4.67 — Shell-script target-shell-parseability mandate (User mandate, 2026-05-19)

Forensic anchor — direct user mandate (verbatim, 2026-05-19):

“any issue we spot must be fixed, bash scripts as well if they are broken!”
 “Make sure that this is mandatory rule!”

Forensic incident — Phase 39.FJ.67 D3 post-flash cycle block: device/rockchip/rk3588/tests/test_all_fixes.sh:114 invoked on Android via sh script.sh used bash-only process substitution `exec > >(tee -a "$file") 2>&1`. Android’s /system/bin/sh is mksh, which parses the *entire* script BEFORE executing it — so the runtime guard `if [ -n "${BASH_VERSION:-}" ]` could not save the script: mksh rejected `>(...` at parse time and the test cycle failed to launch on D3. The runtime guard correctly tried to gate the unsafe path, but mksh’s compile-time parser made the guard useless. Fixed by wrapping the offending `exec` in `eval` so the parser sees only a string and the parse step defers to runtime where the guard can take effect.

This is the canonical class of script-bug FAIL-bluff: a test cycle **PASSes** on the developer host (bash) and **FAILs** at parse time on the target (mksh), producing a FAIL-bluff that masks every product defect the cycle would have caught. §11.4.1 already forbids script-bug FAILs. §11.4.67 mechanically prevents this specific subclass.

The mandate. Every shell script that may be invoked under a target shell other than the one in its shebang **MUST** parse cleanly under that target shell. Specifically:

1. **Closed-set scope.** Every tracked `.sh` file under `device/rockchip/rk3588/tests/`, `scripts/`, and `scripts/testing/` (and equivalent paths in owned submodules) is in scope. Explicitly OUT of scope: `external/`, `prebuilts/`, `packages/modules/`, `kernel-5.10/`, `out/`, `build/`, `scripts/legacy/` (legacy quarantine), and any third-party vendored shell directory the project does not maintain.
2. **Mandatory parseability invariant.** Every in-scope script **MUST** parse under POSIX `sh -n` on the developer host. POSIX `sh` is the closest portable parser to Android `mksh` — passing `sh -n` on the host gives high confidence the script parses on `mksh`.
3. **Bash-only syntax requires deferral.** Process substitution `>(…) / <(…)`, `[ [ ]`, `<<<` here-strings, indexed arrays `arr=()`, associative arrays `declare -A, ${var^^} / ${var,,}` case-fold expansions, `${var/pattern/replacement}` with bash-only operators, `coproc`, function name `{ … }` (vs POSIX `name() { … }`), and any other bash-only construct **MUST** EITHER be wrapped in `eval '…bash-only string…'` so the parser sees only a string OR be guarded by sourcing the file only on bash- detected hosts (`[ -n "${BASH_VERSION:-}" ]` BEFORE the script loads — runtime guards inside a single `mksh`-parsed script do not work).
4. **Honest shebangs.** A script that contains bash-only constructs without `eval`-deferral **MUST** declare `#!/bin/bash` (or equivalent) AND its callers **MUST** invoke it as `bash script.sh` rather than `sh script.sh`. A script with `#!/system/bin/sh` or `#!/bin/sh` shebang **MUST** be POSIX-clean — no bash-only syntax, no `local` keyword without an `mksh` fallback, no `read -p`, no `echo -e` without portable equivalent.
5. **eval-deferral is the safe pattern.** When a bash-only construct is genuinely needed inside a script that may be parsed by `mksh` (e.g. `sh script.sh` from a callsite outside our control), wrap the construct in single-quoted `eval`: `eval 'exec > >(tee -a "$f") 2>&1'`. The parser sees a string; the runtime guard around the `eval` decides whether to execute.

Captured-evidence enforcement. Pre-build gate `CM-SCRIPT-TARGET-SHELL-PARSEABLE` walks every `.sh` file under the in-scope directories, runs `sh -n` on each with a 30-second per-file timeout, and FAILs if any script fails to parse. The gate's diagnostic output cites the failing file + the `sh -n` error message so the fix location is unambiguous.

Propagation gate. `CM-COVENANT-114-67-PROPAGATION` verifies the §11.4.67 anchor literal is present across the 44-file consumer fleet (parent `CLAUDE.md` / `AGENTS.md` + Containers + 10 owned atmosphere submodules + nested SmartTube SharedModules / MediaServiceCore / MSC-SharedModules + 7 HelixQA submodules). Constitution submodule files (`constitution/{Constitution,CLAUDE,AGENTS}.md`) are canonical authority — they carry §11.4.67 by definition.

Paired meta-test mutations. (a) Inject a bash-only construct (`exec > >(/bin/cat)`) outside `eval` into a fresh location inside a test script and assert `CM-SCRIPT-TARGET-SHELL-PARSEABLE` FAILs. (b) Strip 11.4.67 literal from one consumer file and assert `CM-COVENANT-114-67-PROPAGATION` FAILs.

Composes with:

- §11.4.1 (FAIL-bluffs equally forbidden — script-bug FAILs that mask product defects are the exact failure mode §11.4.67 mechanically closes)
- §11.4.4 (test-interrupt-on-discovery — a parse FAIL at cycle start IS the discovery event; cycle stops, fix lands, retest)
- §11.4.6 (no-guessing — `sh -n` produces fact, not opinion)
- §11.4.50 (deterministic consistency — a script that parses on one shell and not another fails the N-iteration identical-result requirement)
- §11.4.51 (live-ADB-first — shell-script fixes are `LIVE_ADB_TESTABLE`: push the fix, re-run on device, capture evidence before commit)

Mandatory protections (no escape hatch):

- **No `--skip-parseability-check` flag.** A script that does not parse under the target shell is broken regardless of intent.
- **No “this script only runs on bash” exception.** If a script may be invoked via `sh script.sh` (and almost all of them can — Android’s default `sh` is `mksh`), it **MUST** parse under `mksh`.
- **No “the runtime guard catches it” exception.** Runtime guards inside a `mksh`-parsed script cannot prevent `mksh` parse-time rejection. Use `eval-deferral` or `shebang-controlled` invocation.
- **Fix at source, not in callsites.** A broken script is fixed in the script itself (per §11.4.1), not by individual callsites defensively detecting parse failures.

Pre-build gate `CM-SCRIPT-TARGET-SHELL-PARSEABLE` + paired mutation.

Propagation gate `CM-COVENANT-114-67-PROPAGATION` + paired mutation.

Canonical authority: constitution submodule [Constitution.md](#) §11.4.67.

Non-compliance is a release blocker regardless of context.

§11.4.68 — Positive sink-side / downstream evidence mandate (User mandate, 2026-05-20)

Forensic anchor — direct user mandate (verbatim, 2026-05-20):

“We still do not hear any audio played from D3 device! Arvus Web Dashboard when we play music from D3 shows nothing for Codec In Use! This **MUST BE** investigated and fixed! How come we passed the tests with Arvus validation? What were values for the Codec In Use field? Empty means nothing! This is not working! It **MUST BE FIXED, TESTED AND VERIFIED WITH FULL AUTOMATION TESTING ASAP!!!**”

§11.4.13 already mandates consuming the Arvus sink-side report when available.

§11.4.68 extends this with a stricter contract: **a test that asserts audio or video routing PASS MUST capture and verify positive sink-side or downstream evidence — never config-only, never metadata-only, never PCM-open-state-only.** The §11.4.13 path described what **SHOULD** be captured; §11.4.68 fixes the failure mode where tests “**PASS**ed” because the helpers silently **SKIP**ped the very evidence the assertion depended on.

Positive sink-side / downstream evidence — closed enumeration (at least one MUST be captured for every audio/video routing PASS; none alone is sufficient if the canonical sink/downstream is reachable):

1. **Sink-side codec-state** (Arvus / Sonos / Yamaha / Denon / Marantz / WiSA REST API) showing a **non-empty Codec-In-Use** value matching the test's expected codec regex during the playback window. Empty / <unreachable> / <N.E.> / <None> placeholders are NOT positive evidence.
2. **PCM frames-written delta** from /proc/asound/cardN/pcmMp/sub0/status hw_ptr — the delta over the playback window MUST be strictly positive AND consistent with the expected sample-rate × channel-count × duration. Zero or negative delta is FAIL.
3. **ALSA ELD / EDID-Like-Data** from /proc/asound/cardN/eld#X.Y demonstrating the sink negotiated the expected channel count + format. Empty ELD is NOT evidence the sink received audio — it only proves the cable / EDID-channel is up.
4. **ffprobe-on-captured-mp4** for video routing — non-zero frame count, expected codec, expected resolution, expected fps. 0-byte capture (Bug #24 pattern) is FAIL.
5. **Recording-analyzer event match** per §11.4.2 / §11.4.5 timeline — at least one matched event in the analyzer findings.
6. **Tinycap RMS amplitude** for ALSA-loopback audio recordings — non-trivial RMS in the captured WAV (≥ -60 dBFS for line-level playback).

Failure-mode taxonomy — what §11.4.68 specifically forbids:

- arvus_probe_present returns 1 → test reports SKIP → suite reports all-green. **FORBIDDEN.** When the sink-side evidence is REQUIRED by the assertion, missing sink is OPERATOR-BLOCKED (an explicit release-blocker, NOT a SKIP, NOT a PASS).
- arvus_probe_codec_state returns empty string → test logs “codec_state:” (blank) → assertion proceeds on HAL-side metadata. **FORBIDDEN.** Empty codec-state IS a defect signal — either the audio is not arriving at the sink (genuine product defect per current User mandate) OR the sink is unreachable (operator-blocked). Either way, the test cannot PASS.
- Test reads HAL dumpsys media.audio_flinger output-thread device field showing 0x400 (AUDIO_DEVICE_OUT_HDMI) → reports “HDMI routing PASS”. **FORBIDDEN.** That field reflects what the audio POLICY chose, NOT what the audio HAL physically opened. The HDMI ALSA card may still be closed (PCM open failure swallowed by the HAL's ALOGW fallback) — silent silence at the user's ear. Sink-side codec-state or PCM hw_ptr delta is the ONLY proof audio actually arrived.
- Test asserts /vendor/etc/audio_policy_configuration.xml contains the deep_buffer route to HDMI Out. **FORBIDDEN as the only evidence.** Config-only PASS is the canonical §11.4 PASS-bluff — config can be perfect AND audio still inaudible (e.g., the HAL may still misroute). MUST be paired with at least one runtime positive evidence from the enumeration above.

Mandatory protections (all four):

1. **Library contract.** Sink-side helper libraries (arvus_probe.sh, sonos_probe.sh, etc.) MUST expose a *_require_reachable function that returns exit code 2 (OPERATOR-BLOCKED) when the sink cannot be probed. Tests asserting sink-side codec-state MUST call this function at probe entry. Silent skip via *_probe_present is permitted ONLY for topology-dispatch

fan-out logic that already has an alternative on-SoC + downstream-amplitude probe path.

2. **Exit-code propagation.** `arvus_require_reachable` / equivalent helpers return 2; the test harness MUST propagate 2 to the suite summary (counted separately from FAIL/SKIP) as OPERATOR-BLOCKED. `test_all_fixes.sh` orchestrator MUST exit non-zero when any test reported OPERATOR-BLOCKED AND release tags MUST NOT be cut while any OPERATOR-BLOCKED test exists.
3. **Test rewrite.** Every test currently asserting audio output routing PASS via metadata-only checks (e.g., the historical `test_audio_output_routing.sh` config-only invariants) MUST be rewritten to capture at least one positive sink/downstream evidence per the enumeration above. The §11.4 PASS-bluff pattern `tinymix says route is on` → PASS is the failure mode this anchor closes.
4. **Anti-stickiness post-stop.** After the test stops the playback, it MUST re-probe the sink/downstream and assert the codec-state transitioned to N.E. / N/A / 0-frames-delta. A persistent codec- state across test boundaries indicates orphan playback (§11.4.14 violation) AND would falsely PASS the next test.

Pre-build gate CM-COVENANT-114-68-PROPAGATION enforces this anchor literal across canonical files + per-consumer propagation. Paired mutation strips the anchor literal → gate FAILs. Pre-build gate CM-POSITIVE-SINK-EVIDENCE-PER-AUDIO-TEST walks every audio-routing on-device test asserting PASS and verifies it cites at least one positive evidence call (one of `arvus_require_reachable` / `arvus_assert_codec_format` / `pcm_hw_ptr_delta_positive` / `tinycap_rms_above_floor` / `ffprobe_frames_nonzero`). Paired mutation strips one such call → gate FAILs.

Composes with §11.4.2 (recorded-evidence — sink-side is the captured-evidence channel), §11.4.5 (audio + video quality analysis comprehensiveness — codec / channel-count / RMS), §11.4.13 (sink-side captured-evidence mandate — §11.4.68 closes its silent-skip gap), §11.4.14 (test cleanup — anti-stickiness post-stop), §11.4.46 (recent-work validation — sink-side evidence required at validation entry), §11.4.49 (dual-approach testing — both UI + Intent variants need sink-side evidence), §11.4.50 (deterministic consistency — sink-side evidence must be consistent across N iterations), §11.4.52 (autonomous-validation — sink probe IS the canonical autonomous path for HDMI audio).

No escape hatch — no `--skip-sink-evidence`, `--allow-empty-codec`, `--sink-unreachable-is-pass`, `--metadata-only-suffices` flag exists anywhere. The discipline exists because the operator/end-user forensic on 2026-05-20 demonstrated the exact failure this anchor forbids: tests reported audio-routing PASS while the user heard nothing and Arvus Codec-In-Use was empty.

Canonical authority: constitution submodule [Constitution.md](#) §11.4.68.

Non-compliance is a release blocker regardless of context.

§11.4.69 — Universal Sink-Side Positive-Evidence Taxonomy + Mechanical Enforcement (User mandate, 2026-05-20)

Forensic anchor — direct user mandate (verbatim, 2026-05-20):

“THIS MUST HAPPEN NEVER AGAIN!!! We MUST HAVE this all working! Not just for audio but for every single piece of the System!!! Proper full automation when executed with success MUST MEAN that manual testing will be as much positive at least regarding the success results! Dive deep into the root causes and how this had happened! The root causes of such omission MUST BE TRACKED, and proper solution applied! Solution MUST BE universal, generic that solves working flows for all System components and for all future and all existing projects! Make sure everything is added we change as mandatory rules we must follow and critical constraints into ours root (constitution Submodule) Constitution.md, CLAUDE.md and AGENTS.md!!! Everything we do MUST BE validated and verified with rock-solid proofs and anti-bluff policy enforcement and fulfillment! THIS IS MOST CRITICAL POINT WE HAVE NOW with all audio issues!”

Escalation from the same operator session (2026-05-20, hours earlier):

“We still do not hear any audio played from D3 device! Arvus Web Dashboard when we play music from D3 shows nothing for Codec In Use! ... How come we passed the tests with Arvus validation? What were values for the Codec In Use field? Empty means nothing! This is not working!”

Forensic incident — 2026-05-19→20 D3 audio-routing PASS-bluff generalised.

§11.4.68 closed the audio-specific failure path where Arvus reported an empty Codec-In-Use field and the test PASSED anyway. The same shape of bug had previously shown up intermittently across multiple subsystems (display routing on D1 secondary HDMI, BT A2DP codec advertisements, WiFi connection quality, video playback frame-count, touch-input event delivery). The audio incident was the canonical trigger because the operator hears silence at the soundbar even while the cycle reports green. §11.4.69 is the universal generalisation — closing the bluff class for every user-visible feature, not only audio.

Root-cause analysis — why existing anchors §11.4.2 / §11.4.5 / §11.4.13 / §11.4.27 / §11.4.52 / §11.4.68 did NOT mechanically catch the failure across the System:

1. **Aspirational text without a pre-build gate.** §11.4.2 (recorded-evidence) required captured visual/audio evidence per PASS but had no gate walking the test corpus to enforce a per-PASS evidence file. Tests called `ab_pass` "description" and the evidence requirement was honour-system.
2. **SKIP-fail-open hiding missing evidence.** §11.4.13 (Arvus sink-side) had a silent-skip escape when the sink returned an empty / unreachable response — the empty response was the very defect signal but it was classified as “sink unreachable, SKIP the assertion.” §11.4.68 closed this for audio; §11.4.69 closes it universally across every feature class with a downstream sink.
3. **PASS-by-default vs FAIL-by-default helper contract.** The bare `ab_pass` helper accepted a free-text description and incremented the PASS counter — no evidence path required, no shape enforced. The default failure mode was therefore PASS- bluff. The correct default is FAIL-by-default — a PASS must prove itself with a captured-evidence artefact path.
4. **Missing closed-set taxonomy.** Individual test authors made one-off decisions about what evidence shape was acceptable. Without a single canonical table mapping every user-visible feature class to its required sink-side probe, aggregate enforcement across the System was mechanically impossible.

5. Helper functions that mapped a multi-step probe to one bool.

arvus_probe_present collapsed reachable-but-empty into the same value as unreachable. The §11.4.68 fix added arvus_require_reachable for audio; §11.4.69 codifies the equivalent for every sink (Sonos, Yamaha, Denon, video display analyzer, WiFi sink, BT peer, etc.).

6. No per-feature evidence ledger. §11.4.52 supplied an autonomous-validation classification (AUTONOMOUS_VERIFIED etc.) but did not require each test to cite a captured-evidence artefact path matching the §11.4.69 taxonomy.

The §11.4.69 mandate — universal, generic, applies to every System component, every owned project, every consumer subscribing via the constitution submodule:

Element 1: Closed-set sink-side / downstream evidence taxonomy. Every user-visible feature class in every consuming project **MUST** map to exactly one entry in the following table. The taxonomy is the canonical authority on the required evidence shape per feature class. Tests annotate themselves with # §11.4.69 FEATURE: <class> so the pre-build gate can match expected evidence shape.

Feature class	Required probe	Required evidence shape
audio_output	sink-side codec/channel REST probe (Arvus / Sonos / Yamaha / Denon / Marantz / WiSA) + on-device tinycap capture + ffprobe channel assertion	non-empty codec_in_use matching expected codec AND captured WAV with RMS > -60 dBFS AND ffprobe channels matches claim
audio_input	tinycap capture from intended device + ffprobe RMS	captured WAV with RMS > -60 dBFS AND audio_devices_for_attr shows the claimed source device
video_display	screenrecord on intended display + ffprobe -count_frames + analyzer event-match	non-zero mp4 size AND nb_read_frames > 0 AND analyzer reports a matched event for the claimed display
network_throughput	iperf3 / curl --speed / dd over network	throughput ≥ per-link-type floor AND TCP buffer / congestion-control matches claim
network_connectivity	ping + DNS resolve + TCP connect probe	captured ping reply AND DNS answer AND TCP handshake to known port
bluetooth_a2dp	codec query + sink peer state + A2DP TX queue depth	dumpsys bluetooth_manager reports CONNECTED+A2DP_PLAYING AND codec field non-empty AND queue depth < overflow threshold
bluetooth_pair	dumpsys bluetooth_manager + BluetoothDevice.getBondState()	

Feature class	Required probe	Required evidence shape
		bond state = BOND_BONDED AND device MAC matches expected
touch_input	getevent -lt capture during scripted input tap	≥ 1 EV_ABS ABS_MT_* event per scripted tap AND dumpsys input shows the touch device active
sensor	dumpsys sensorservice + getevent on sensor device	event count > 0 between BEFORE and AFTER snapshots for the claimed sensor
gpu_render	SurfaceFlinger frame-rate dump + dumpsys gfxinfo <pkg>	Janky frames < threshold AND Total frames rendered > 0 during test window
storage_read	dd if=<path> of=/dev/null + / proc/diskstats delta	throughput $\geq$ floor AND read_sectors delta confirms IO landed on intended device
storage_write	dd of=<path> + sync + /proc/ diskstats delta	throughput $\geq$ floor AND write_sectors delta confirms IO landed on intended device
mediacodec_decode	media.metrics dump + dumpsys media.codec	decoder lifecycle events present (configure, start, ≥ 1 output buffer) AND codec name matches claim
mediacodec_encode	media.metrics dump + non-zero output file	encoder lifecycle events present AND output file non-zero AND ffprobe confirms codec matches claim
miracast / cast	sink-side WFD / Cast announce + session state	sink dashboard reports active session AND session ID present in dumpsys media_session
boot_service	getprop init.svc.<name> + log capture	service state = running AND log shows expected boot-time output literal
package_install	pm path <pkg> + signature match	non-empty pm path AND APK signature matches expected platform key
permission_grant	dumpsys package <pkg> permission state	permission state = granted: true for the claimed permission
wifi_link	iw dev wlan0 link + RSSI + frequency	captured BSSID AND RSSI within range AND frequency matches claim

Feature class	Required probe	Required evidence shape
wifi_throughput	iperf3 to LAN server	throughput $\geq$ floor for the negotiated PHY rate
ethernet_link	ip link show + carrier state + throughput probe	carrier=1 AND speed matches claim AND throughput probe succeeds
display_topology	DRM connector dump + display ID enumeration	connector status connected for claimed connector AND display ID present in dumpsys SurfaceFlinger
drm_playback	Widevine session + decoded-frame count	active DRM session AND decoded-frame count > 0
subtitle_render	analyzer + Tesseract OCR on captured frames	OCR matches expected subtitle text on the claimed display

The taxonomy is **open to additions** via constitution-submodule commits + propagation. Removal of a feature class is forbidden without a tracked work item explaining the supersession. Consumer projects MAY extend the taxonomy with project-specific feature classes (e.g., HelixCode-specific build-pipeline classes) but MUST NOT contract it.

Element 2: New helper contracts (additive during grace period; mandatory after grace-period end-date 2026-06-19).

The project anti-bluff helper library (anti_bluff.sh in ATMOSphere; functionally equivalent helper in every consuming project) MUST provide three new helpers:

1. **ab_pass_with_evidence** <description> <evidence_path> — the new canonical PASS helper:
 - REQUIRES a non-empty evidence_path argument
 - Verifies the path exists AND is non-empty ([-s "\$path"])
 - If the path is missing or empty, the helper EXITs the test as FAILED with a diagnostic citing the missing evidence path
 - On success, increments the PASS counter AND emits a result line PASS: <description> [evidence: <path>]
 - Composes with ab_send_action — a single test issues many ab_send_action calls and one terminal ab_pass_with_evidence
2. **ab_skip_with_reason** <description> <reason_code> — the new canonical SKIP helper:
 - REQUIRES reason_code from a CLOSED set: geo_restricted, operator_attended, hardware_not_present, topology_unsupported, network_unreachable_external, feature_disabled_by_config
 - REJECTs ad-hoc reasons — any string not in the closed set fails the helper at call time
 - FORBIDS network_unreachable_external for any feature class in the §11.4.69 taxonomy that lists a sink-side probe — that escape was the §11.4.13 fail-open path §11.4.68 closed for audio, and §11.4.69 closes universally
 - SKIP remains mechanically distinct from PASS in cycle math (per §11.4.1)

3. **Bare ab_pass deprecation wrapper.** The legacy ab_pass helper MUST emit a WARNING line on every call AND emit a deprecation event to a tracked log file (/data/local/tmp/ab_pass_deprecations.log). After the grace period (defined in Element 4) the wrapper MUST FAIL the test outright, forcing all callers to migrate.

Element 3: Mechanical enforcement via three pre-build gates + three paired meta-test mutations (per §1.1).

1. **CM-SINK-EVIDENCE-PER-FEATURE** — walks every .sh test in device/rockchip/rk3588/tests/ (and equivalent project paths), parses for # §11.4.69 FEATURE: <class> annotation comments, verifies the test invokes the corresponding sink-side probe from the taxonomy AND uses ab_pass_with_evidence (or ab_skip_with_reason with a permitted reason). Pre-grace: tests lacking the annotation WARN. Post-grace: they FAIL.
2. **CM-NO-FAIL-OPEN-SKIP** — audits sink-side probe helpers (arvus_probe.sh, sonos_probe.sh, equivalents) for any code path that converts an empty / unreachable sink response into a PASS-counting SKIP for a feature class with a sink-side probe in the taxonomy. The gate FAILs if any such path exists. Remediation: escalate to FAIL (the empty response IS the defect) or to OPERATOR-BLOCKED per §11.4.68.
3. **CM-AB-PASS-WITH-EVIDENCE-EVERYWHERE** — pre-grace WARNs on every bare ab_pass call in in-scope tests; post-grace (2026-06-19) FAILs the gate. Tests that legitimately have no evidence-capable PASS path MUST migrate to ab_skip_with_reason operator_attended and add a tracked work item per §11.4.52's OPERATOR_ATTENDED_ONLY classification.

Paired mutations (per §1.1):

- **Mutation 1:** strip the literal ab_pass_with_evidence from the project anti-bluff library → assert CM-AB-PASS-WITH-EVIDENCE-EVERYWHERE FAILs.
- **Mutation 2:** inject a fresh test that calls bare ab_pass without the # §11.4.69 FEATURE: annotation → assert CM-SINK-EVIDENCE-PER-FEATURE FAILs post-grace, WARNs pre-grace.
- **Mutation 3:** inject a fresh test that calls ab_skip (no _with_reason suffix) OR ab_skip_with_reason network_unreachable_external for an audio_output feature → assert CM-NO-FAIL-OPEN-SKIP FAILs.

Element 4: Grace period mechanics — additive, not destructive.

- **Grace period:** 30 days from anchor land date — 2026-05-20 through 2026-06-19 inclusive. The end date is hardcoded in the helper library and in this anchor; no environment variable, no command-line flag, no operator override can extend or shorten it.
- **Pre-grace behaviour:** new helpers exist and are usable; the three gates emit WARN for legacy patterns; cycle math counts WARNs separately from PASS / FAIL / SKIP.
- **Post-grace behaviour:** WARN promotion to FAIL is automatic via a date check (["\$(date -u +%Y%m%d)" -ge 20260619]) inside each gate's enforcement logic.

Element 5: Inheritance via HelixConstitution. This anchor lives in the constitution submodule and applies to every consuming project that subscribes — including HelixCode, Catalogizer, Yole, HelixPlay, HelixTranslate, HelixFlow, MeTube, ATMOSphere-Android-15, and every future consumer. Each consumer

MUST implement its own equivalent helper library + pre-build gates; the taxonomy in Element 1 is the canonical authority across all consumers. Consumer-specific extensions are allowed only as additions to the taxonomy.

Composes with:

- §11.4.1 (FAIL-bluffs equally forbidden — `ab_skip_with_reason` closed-set prevents script-bug FAILs hiding real defects)
- §11.4.2 (recorded-evidence requirement — Element 2 helper contracts make captured evidence mechanically required per PASS)
- §11.4.5 (audio + video quality 5-layer — §11.4.69 enforces that the layers are exercised before PASS via the taxonomy probe)
- §11.4.6 (no-guessing — sink-side evidence eliminates “we think audio is routing” guessing; the codec field reads concretely)
- §11.4.13 (Arvus sink-side evidence — §11.4.69 mechanically closes the SKIP-fail-open escape §11.4.13 left open)
- §11.4.27 (no-fakes-beyond-unit — §11.4.69 adds the runtime hook §11.4.27 lacked)
- §11.4.50 (deterministic consistency — sink-side evidence is what `ab_run_n_times` hashes for divergence detection)
- §11.4.52 (autonomous-validation — §11.4.69 supplies the evidence-shape contract per feature class; §11.4.52 supplies the classification ledger)
- §11.4.68 (positive sink-side / downstream evidence for audio + video — §11.4.69 is the universal generalisation across all feature classes)

No escape hatch:

- No `--skip-evidence` flag.
- No `--config-only-pass` flag (configuration-only PASS is the exact bluff pattern this anchor closes).
- No `--allow-fail-open-skip` flag.
- No `--legacy-ab-pass-permitted` flag post-grace.
- No taxonomy-bypass shortcut — every feature class in scope MUST have its evidence shape enforced.

Propagation gate. CM-COVENANT-114-69-PROPAGATION verifies the §11.4.69 anchor literal is present across the ~44-file consumer fleet (parent CLAUDE.md / AGENTS.md + Containers + 10 owned atmosphere submodules including the smarttube-player nested set + 7 HelixQA submodules). Constitution submodule files (constitution/{Constitution,CLAUDE,AGENTS}.md) are canonical authority — they carry §11.4.69 by definition.

Canonical authority: constitution submodule Constitution.md §11.4.69.

Non-compliance is a release blocker regardless of context.

§11.4.70 — Subagent-Driven Execution Is The Default (User mandate, 2026-05-20)

Forensic anchor — direct user mandate (verbatim, 2026-05-20):

“Always do if possible Subagent-driven! Add this into our root (constitution Submodule) Constitution.md, CLAUDE.md and AGENTS.md. This should be the default choice ALWAYS!”

When executing implementation plans authored via `superpowers:writing-plans` (or any equivalent task-decomposed execution flow), the **default execution model is subagent-driven** per `superpowers:subagent-driven-development`. Inline execution via `superpowers:executing-plans` is permitted ONLY when (a) the task is trivial AND fits in a single sub-300-line edit, OR (b) the operator explicitly requests inline execution at brainstorm-handoff time per the `superpowers:writing-plans` skill's explicit-execution-mode prompt.

Why subagent-driven is the default:

1. **Isolated context per task** — each subagent gets a fresh context window scoped to its task, avoiding the conductor's context bloat that degrades reasoning quality after ~100 k tokens.
2. **Two-stage review naturally enforced** — fresh subagent does the work, conductor reviews, optional second subagent verifies. This composes with §11.4.4 four-layer coverage + §11.4.43 TDD-fix- discipline.
3. **Parallel-PWU compatible** (§11.4.58) — subagents are the unit of parallelism in the PWU pipeline; making them the default aligns the writing-plans handoff with the parallel-development methodology.
4. **Anti-bluff seam** (§11.4) — the conductor's review of a subagent's output is structurally separated from the work itself, eliminating self-review blind spots.
5. **Survives operator absence** — subagents resume from their on-disk plan + spec inputs; the conductor session can be lost without losing in-flight task state.

Composition with §11.4.4 (four-layer coverage), §11.4.6 (no- guessing — subagent's captured output IS the evidence), §11.4.42 (iteration discipline), §11.4.43 (TDD-fix), §11.4.50 (deterministic consistency — the subagent's deterministic exit code), §11.4.51 (LIVE_ADB_FIRST — subagents classify rebuild-requirement before commit), §11.4.58 (parallel-development PWU — subagents ARE the parallel work units).

No escape hatch. `--inline-execution-required`, `--no-subagents`, `--monolithic-execution` are NOT permitted flags. Operator may request inline at brainstorm-handoff time per `superpowers:writing-plans` skill; absent that, subagent-driven is default. Skipping subagent-driven for non-trivial work without recorded operator authorisation is itself a §11.4 PASS-bluff (the conductor self-reviews, the anti-bluff seam disappears, and the forensic class of incident this Constitution catalogues becomes mechanically possible again).

Captured-evidence enforcement. Pre-build gate `CM-COVENANT-114-70-SUBAGENT-DEFAULT-PROPAGATION` enforces this anchor literal in every `CLAUDE.md` / `AGENTS.md` across parent + 10 owned submodules + nested submodules + HelixQA dependencies (the ~44-file consumer fleet). Paired §1.1 meta-test mutation strips the anchor literal → gate FAILs.

Canonical authority: constitution submodule `Constitution.md` §11.4.70. Default applies to every `superpowers:writing-plans` skill terminal handoff and every multi-task execution flow with independent subtasks.

Non-compliance is a release blocker regardless of context.

§11.4.71 — Pre-Push Fetch + Investigate + Integrate Mandate (User mandate, 2026-05-20)

Forensic anchor — direct user mandate (verbatim, 2026-05-20):

“before pushing changes to any upstream for any repository - main repo or Submodule, we MUST fetch and pull all changes. Once these are obtained WE MUST investigate what is different compared to head position we were on last time before fetching and pulling new changes! We MUST understand what is done and for what purpose, easpecially how that does affect our project and our System in general! Any mandatory changes or improvements required by fresh changes we just have brought in MUST BE incorporated, covered with all supported types of the tests which will produce as a result of its success execution REAL PROOFS of working for all componetns and functionalities covered and work fully in anti-bluff manner!”

This anchor is the **everyday-push variant** of §11.4.41 (Pre-Force-Push Merge-First). §11.4.41 governs the destructive force-push case; this anchor governs EVERY push including ordinary fast-forwards. The complementary discipline closes the regression-via-stale-baseline gap: a project committing against a HEAD that's behind any upstream's main risks pushing changes that conflict with, regress, or fail to incorporate accepted work landed in parallel by other consumers (HelixCode, Catalogizer, Yole, HelixPlay, HelixTranslate, HelixFlow, MeTube on the constitution submodule; sibling projects on other shared submodules).

The mandatory 5-step pre-push cycle (every push, every repository — main + every submodule):

1. **Fetch all remotes** — `git fetch --all --prune --tags` (or per-remote loop if `--all` is unavailable). Capture stdout for the audit trail.
2. **Pull all upstream branches** — `git pull --no-rebase <remote> <branch>` for each configured remote whose tip differs from the local branch. Resolve any merge conflict per §11.4.41 step 2 (rebase / merge / cherry-pick per consumer judgment, NOT auto---ours or --theirs).
3. **Investigate the diff vs OUR previous HEAD** — `git log --oneline <previous-head>..HEAD + git diff --stat <previous-head>..HEAD` + read EVERY commit's body. For each foreign commit understand: (a) what changed, (b) why (forensic anchor + cited user mandate or §-anchor), (c) how it affects OUR project + OUR System (does it touch surfaces we ship? does it introduce constraints we now MUST honor? does it deprecate or supersede patterns we use?).
4. **Integrate mandatory changes + cover with full anti-bluff test coverage** — if the pulled-in work mandates anything (new §-anchor, new gate, new propagation requirement), implement it per §11.4.4(b) four-layer coverage (pre-build gate + post-build inspection + on-device test + HelixQA Challenge) + §11.4.43 TDD-fix RED-test-first discipline. Every PASS MUST carry §11.4.5 captured-evidence (REAL PROOFS — not metadata-only).
5. **Then push** — only after Steps 1–4 land. Push to every configured remote in the per-repo cascade order (canonical first, then mirrors). Verify the push landed with `git ls-remote <remote> <branch>` post-push.

Composition with §11.4.26 (constitution-submodule update pipeline — per-submodule specialisation) + §11.4.32 (post-pull validation) + §11.4.37 (fetch-before-edit) + §11.4.40 (full-suite retest before tag) + §11.4.41 (pre-force-push

merge-first — the force-push case of this anchor) + §11.4.42 (iteration discipline) + §11.4.43 (TDD-fix-discipline) + §11.4.4(b) (four-layer coverage) + §11.4.5 (audio + video quality analysis comprehensiveness) + §11.4.6 (no-guessing — every integration decision cites the foreign commit by SHA, not “we think it does X”).

No escape hatch — there is no `--skip-fetch`, `--no-investigate`, `--fast-push`, `--trust-upstream` flag. The discipline exists because:

- Pushing-without-fetching produces silent stale-baseline regressions (consumer A’s PR diverges from consumer B’s PR; neither saw the other; both merge; resulting tree corrupts).
- Pulling-without-investigating destroys the foreign work’s traceability (we now carry someone else’s code without understanding it; future debugging crosses the §11.4.6 no-guessing line).
- Integrating-without-testing reproduces the §11.4 PASS-bluff pattern (we adopt a change but don’t prove it works in OUR system).
- Pushing-after-integration-without-evidence reproduces §11.4 itself (untested integration = release-blocker by definition).

Per-repository scope: this anchor applies to (a) the parent ATMOSphere-Android-15 repo, (b) the constitution submodule (across all 6 remotes), (c) every owned submodule (presenter, vlc-player, nova-player, mpv-player, gramophone-player, rhythm-player, strep-player, smarttube-player, torrservice, lampa), (d) every nested submodule (smarttube-player/SharedModules / smarttube-player/MediaServiceCore / smarttube-player/MediaServiceCore/SharedModules), (e) every HelixQA dependency (Challenges, Containers, DocProcessor, LLMOrchestrator, LLMProvider, VisionEngine, HelixQA).

Audit-trail requirement: every push event MUST be reconstructable post-hoc from `docs/changelogs/<tag>.md` + the per-repo `git log` evidence. Operators auditing a release MUST be able to: (1) see what each repo’s HEAD `.@{u}` was BEFORE this push, (2) see what foreign work was pulled in, (3) see what new tests + captured-evidence were generated by the integration. Anchor enforcement may be partially automated by extending `scripts/commit_all.sh` to refuse pushes without an audit-trail log entry.

Pre-build gate CM-COVENANT-114-71-PROPAGATION enforces this anchor literal across the consumer fleet (parent + 10 owned submodules + nested + HelixQA). Paired mutation strips the anchor literal → gate FAILs. Paired meta-test also asserts the audit-trail-log file produced by a recent push exists + lists the foreign-commit SHAs that were pulled in.

Canonical authority: constitution submodule [Constitution.md](#) §11.4.71.

Non-compliance is a release blocker regardless of context.

§11.4.72 — Audio Top-Priority Mandate (User mandate, 2026-05-20)

Forensic anchor — direct user mandate (verbatim, 2026-05-20):

“Make sure all fixes for audio are always top priority in main working stream!”

The conductor (main working stream — whether a Claude Code session, an AI agent, or a human operator) **MUST** treat audio fixes as the highest-priority class on the serial dispatch queue. Audio fixes include but are not limited to §EU HDMI rejection, §EM multichannel HDMI, §ET Arvus integration, §EV/§EW/§EX D3 audio defects, HiFi (Fix #74 pcm_config_deep, #76 HDMI/SPDIF resilience), AC3 (Fix #106 FFmpeg Codec2), ES8388 (Fix #103/§104/#105 mutual-output + DAPM), multichannel LPCM (Fix #112), and every future audio-stack improvement.

Composition with §11.4.42 (iteration-discipline mandate) + §11.4.58 (parallel-development methodology). Per §11.4.42, the conductor selects work items in priority order; this anchor pins audio at the apex of that order. Per §11.4.58 PWU pipeline, parallel BACKGROUND subagents (research, refactors, infrastructure documentation) **MAY** run concurrently with audio work — they do **NOT** preempt audio on the main-stream serial dispatch queue.

Operative rule: any time the conductor faces a choice between dispatching an audio task vs a non-audio task on the **SAME** serial resource (most commonly the pre_build_verification.sh + meta_test_false_positive_proof.sh edit-and-commit critical path), the audio task wins.

No escape hatch — there is no “but this non-audio task is faster” or “but this research is more interesting” override. The discipline exists because audio-stack regressions are user-perceptible and high-impact (D3 silent post-flash is a release blocker), while research and refactors can wait.

Canonical authority: constitution submodule Constitution.md §11.4.72.

Non-compliance is a process violation regardless of context.

§11.4.73 — Main-specification document versioning + revision discipline (User mandate, 2026-05-20)

Forensic anchor — direct user mandate (verbatim, 2026-05-20):

“Make sure everything we add now in previous and upcoming requests IS ALWAYS applied to the main specification — if we have one. Since all these are not major changes we could increase Specification version per change for secondary version instead of the primary. Primary version **MUST BE** increased for much bigger levels of changes! Add this into root (constitution Submodule) Constitution.md, CLAUDE.md and AGENTS.md as mandatory rule / constraint applicable **ONLY IF** we have something like the main specification document or we do recognize something like the main specification document. Document **MUST BE** updated **ALWAYS** to follow the versioning rules we are applying here + revision and other properties we have!”

Scope condition. This clause applies **only when a project recognises a main specification document** — i.e. a project-level Markdown file (or set of files) that constitutes the canonical “what this system does” reference, typically at docs/specs/**/specification*.md or a comparable path. Projects with no main specification (small libraries, internal tooling) are exempt. Projects that operate per-feature spec files (e.g. ADRs only) are exempt unless they elect to adopt the main-spec pattern.

The mandate. When a project DOES have a main specification document:

1. **Every additive operator requirement, refinement, or accepted recommendation MUST be applied to the spec** before, or as part of, the work that implements it. The spec is the source of truth; downstream code and tracking docs reference it.
2. **Spec versioning has two axes** — *primary* and *secondary* (where “secondary” maps to the existing §11.4.61 metadata-table Revision row):
 - **Primary (V1 / V2 / V3 / ...)** bumps for major rewrites — substantial architecture changes, foundational scope shifts, deprecating large sections, replacing the technology stack. Triggered explicitly by operator decision; old primary versions move to archive/ per §11.4.61.
 - **Secondary (Revision)** bumps for every other change — section additions, mandate landings, structural reorganisation, polish, type-contract refinements. The metadata table’s Revision integer is the secondary version (matches §11.4.44 monotonic-integer rule).
3. **The metadata table on the spec MUST stay current** — Revision bumps in lockstep with the change; Last modified updates to the change date; Status summary describes the bump’s content; Fixed and Fixed summary reference the relevant work IDs.
4. **Cross-document propagation** — when a project has propagated copies of the spec-change rule (Herald uses CLAUDE.md, AGENTS.md, HERALD_CONSTITUTION.md per its §106), those copies MUST also reference the active spec file (specification.V<primary>.md) and not a stale archived version.
5. **Archive semantics** — when the primary version bumps, the old specification.V<n>.md moves to <spec-dir>/archive/ specification.V<n>.md and gains Status: superseded with a Continuation pointer to the new active file. Per the §11.4.65 universal-export mandate the archived .html and .pdf siblings travel with it.

Anti-bluff captured-evidence gate (planned). CM-SPEC-VERSION-DISCIPLINE: - If the project advertises a main spec (operator marks via [project].main_spec_path = "docs/specs/.../specification.V<n>.md" in a canonical config file), the gate asserts: - The file at main_spec_path exists. - Its metadata table’s Revision row is $\geq$ the highest Revision referenced anywhere else in the repo (Issues.md, Fixed.md, CLAUDE.md, AGENTS.md). - If commits exist that touch the spec since the last Revision bump, the gate FAILs (operator forgot to bump).

Paired §1.1 mutation. Backdate the Last modified row to a prior date while leaving real-content edits in place → gate FAILs.

Composition. Composes with §11.4.44 (revision header — Revision is the secondary version here), §11.4.61 (metadata table + ToC), §11.4.59 (README always-sync — if README cites the spec, the citation MUST point at the current primary version), §11.4.65 (universal export — archived versions stay exported).

Classification: universal (per §11.4.17), applicable conditionally per the scope condition above.

Canonical authority: constitution submodule Constitution.md §11.4.73.

Non-compliance (forgetting to bump, or letting the spec drift behind operator-mandated requirements) is a release blocker.

§11.4.74 — Submodule-catalogue-first discovery + extend-don't-reimplement (User mandate, 2026-05-20)

Forensic anchor — direct user mandate (verbatim, 2026-05-20):

“We MUST ALWAYS check which already developed features / functionalities do exist as a part of our comprehensive Submodules catalogue located in vasic-digital and HelixDevelopment organizations on GitHub and GitLab both! Project MUST BE aware of all its existence so we do not implement same things multiple times if they are already done as some of existing universal, reusable general development purpose Submodules! For any missing features that some Submodules we incorporate may be missing we MUST IMPLEMENT the properly and extend those Submodules further! We do control all of the and we CAN and MUST maintain and extend the regularly! All development cycle rules we have MUST BE applied to them and fully respected!”

The mandate. Before scaffolding ANY new module, package, helper, or utility, the contributor (human or AI agent) MUST:

1. **Survey the canonical Submodule catalogue.** Two organisations are authoritative:
 - vasic-digital on GitHub (<https://github.com/vasic-digital>) AND GitLab (<https://gitlab.com/vasic-digital>).
 - HelixDevelopment on GitHub (<https://github.com/HelixDevelopment>) AND GitLab (<https://gitlab.com/HelixDevelopment>). Both organisations mirror across all four supported hosts per §2.1 — surveying one mirror MUST yield the full catalogue.
2. **Inventory existing Submodules.** The canonical repo INDEX is submodules-catalogue.md in this submodule (landed 2026-05-20 per user follow-up mandate). It enumerates every owned-by-us repository under vasic-digital + HelixDevelopment (142 repos at landing time) with a one-line capability description, categorised so consuming projects can locate “is there already a Submodule for X?” in one glance. The catalogue MUST be re-generated quarterly (or after any non-trivial repo create/rename) per its §3 regeneration recipe.
3. **Reuse before reimplement.** If a Submodule already provides the functionality (or 80%+ of it), the consuming project MUST add the existing Submodule as a Git submodule, not write the functionality fresh.
4. **Extend in-place when 80%+ matches but features are missing.** When an existing Submodule is close-but-not-complete, the missing features MUST be added TO THAT SUBMODULE (PR upstream + bump consuming project’s submodule pointer) — never as a separate consuming-project helper that duplicates the 80% already in the Submodule.
5. **Same development-cycle rules apply.** Every Submodule under the two orgs is subject to: §11.4 anti-bluff covenant, §1.1 paired-mutation tests, §11.4.10 credentials handling, §11.4.61 metadata table + ToC, §11.4.65 universal Markdown export, §11.4.73 spec versioning (if the Submodule has its own spec), §2.1 multi-mirror push, §3 propagation order. Maintenance + extension of those Submodules MUST follow these rules without exception.

6. **Document the survey result.** Each new feature's tracker entry (Issues.md row, ADR, PR description, or equivalent) MUST contain a Catalogue-Check: field with one of three values:

- Catalogue-Check: reuse <org/repo>@<sha> — the existing Submodule covers this.
- Catalogue-Check: extend <org/repo>@<sha> — extending an existing Submodule; reference the upstream PR.
- Catalogue-Check: no-match <date> — surveyed both orgs on the named date; no existing Submodule covers this functionality; the new code is justified to live in-project. Operators reviewing the PR re-validate this claim.

Anti-bluff captured-evidence gate (planned). CM-SUBMODULE-CATALOGUE-CHECK:
- For every new module / package introduced in a non-trivial PR (heuristic: ≥ 200 LOC of new non-test code), the gate scans the PR description and any linked tracker row for a Catalogue-Check: line. Absence is a FAIL. - The gate's paired mutation strips the line from a PR and asserts the gate FAILs.

Composition. Composes with §3 (submodule commits propagate first), §4 (every tag on the main repo MUST be mirrored on every owned submodule), §11.4.36 (mandatory install_upstreams on clone/add), §11.4.31 (Submodule-Dependency-Manifest), §11.4.32 (post-pull validation), §11.4.28 (Submodules-As-Equal-Codebase).

Why this matters. Without the catalogue-check, consuming projects re-implement UUIDv7 helpers, RLS-tenant-context wrappers, Pandoc-multi-format exporters, install_upstreams shell scripts, Markdown ToC generators, fnv1a32 hash wrappers, etc. — fragmenting the catalogue and forcing the next project to choose between three subtly-different implementations of the same thing. The mandate makes the catalogue load-bearing.

Classification: universal (per §11.4.17). Applies to every project that consumes this Constitution.

Canonical authority: constitution submodule Constitution.md §11.4.74.

Non-compliance is a process violation; severe cases (duplicate implementation landed without catalogue check) are release blockers.

§11.4.75 — Mechanical Enforcement Without Exception (User mandate, 2026-05-20)

Forensic anchor — direct user mandate (verbatim, 2026-05-20):

“Why do these violations still happen!? This is a serious problem! We cannot rely on stability nor consistency if we cannot respect our Constitution, mandatory rules and constraints! Is there a way to make this always respected, followed and applied without exception fully and unconditionally!? WE MUST HAVE THIS WORKING FLAWLESSLY!!! Do investigate the root causes of such problems! Once all problems are identified WE MUST apply proper mechanisms for this not to happen NEVER EVER AGAIN!”

§-slot history note. This anchor was originally drafted as §11.4.74 but renumbered to §11.4.75 per the §11.4.41 merge-first mandate after upstream landed a different §11.4.74 (Submodule-catalogue- first discovery) concurrent with this work. Same renumber pattern as Phase 39.DH §11.4.41 collision resolution (commit 8ad51547115 of the parent ATMOSphere-Android-15 repo).

The §11.4 covenant (“tests pass while feature broken-for-end-user”) historically relied on agent and operator vigilance to enforce. Three forensic incidents in 2026-05-19→20 demonstrated the failure: subagent a84d2c86f90fdc297 committed docs/research/workstation/Configurations.md without HTML+PDF siblings (§11.4.65 violation) when it stalled at the 600s watchdog before its pandoc-export step; subagent a6b38eedf4ef796e8 dispatched to remediate ALSO stalled at the same operational seam; the parent CLAUDE.md / AGENTS.md drift after §11.4.66 propagation required a dedicated catch-up commit f93b25a92eb. The common failure mode: late-binding enforcement that fires at pre_build_verification.sh time — hours-to-days after the violator commit reached every remote.

The mandate. Constitutional invariants **MUST** be enforced mechanically at FIVE independent layers — bypassing any single layer does not bypass the discipline. The five layers are:

1. **Local pre-commit git hook** — refuses staged .md lacking sibling .html+.pdf (and other staged-only invariants).
2. **commit_all.sh integration** — the canonical project commit script invokes the same checks + auto-runs sync_all_markdown_exports.sh to self-repair before the commit.
3. **Local pre-push git hook** — re-runs siblings + propagation gate subset on every commit in the push.
4. **post-commit auto-repair hook** — detects orphan .md in the just-committed manifest, auto-generates siblings via pandoc + weasyprint, creates a chore(§11.4.75): auto-export ... follow-up commit. Idempotent + recursion-guarded.
5. **Local-only equivalent** (Phase 39.GF, User mandate 2026-05-20). Remote CI surfaces (GitHub Actions, GitLab pipelines, Jenkins, CircleCI, etc.) are **DISABLED** per User mandate 2026-05-20 — the workflow file is preserved at .github/workflows/constitution-compliance.yml.disabled-local-only so its contents survive for re-enable but GitHub Actions does NOT execute it (only .yaml / .yml files under .github/workflows/ are honoured). Layer 5 enforcement migrated to the LOCAL pre-build- verification + meta-test run that the operator **MUST** execute before tagging: bash device/rockchip/rk3588/tests/pre_build_verification.sh
 - bash scripts/testing/meta_test_false_positive_proof.sh.Layers 1-4 remain authoritative; Layer 5 is now the operator’s local final-gate ritual. A future re-enable PWU may re-establish remote CI.

Helper contracts (mandatory):

- scripts/install_git_hooks.sh — idempotent installer that copies hooks from tracked scripts/git_hooks/ into .git/hooks/. Hooked into scripts/setup.sh so fresh clones auto-install. Verifies itself with --verify flag.
- scripts/git_hooks/pre-commit — Layer 1 implementation.
- scripts/git_hooks/pre-push — Layer 3 implementation.
- scripts/git_hooks/post-commit — Layer 4 implementation.
- scripts/git_hooks/commit-msg — Bypass-rationale enforcement.

- `_constitution_sibling_check` function in `scripts/commit_all.sh` — Layer 2 implementation.

Bypass policy. No layer is perfectly bypass-proof in isolation — that is the design. Bypassing ALL FIVE simultaneously is mechanically impossible because:

- Layers 1 + 3 can be bypassed by `--no-verify`. Layer 5 (now local pre-tag ritual after Phase 39.GF) re-runs the same check before tag creation; tags are NEVER created without it per §11.4.40.
- Layer 2 can be bypassed by direct `git commit`. Layer 1 catches that.
- Layer 4 cannot be bypassed by `--no-verify` (runs after commit lands).
- Layer 5 is now local-only (Phase 39.GF — User mandate 2026-05-20). The operator runs `pre_build_verification.sh + meta-test` before every tag per §11.4.40 full-suite-retest mandate. Skipping the local Layer 5 ritual blocks tag creation.

Legitimate bypasses MUST be audited. When `--no-verify` is detected via touch-file marker `.git/ATMO_LAST_BYPASS_ATTEMPT`, the `commit-msg` hook REQUIRES a Bypass-rationale: `<reason>` footer in the commit message. `docs/audit/bypass_events.md` accumulates the audit trail per §11.4 captured-evidence requirement.

Captured-evidence enforcement. Five pre-build gates with paired §1.1 meta-test mutations:

- `CM-COVENANT-114-75-PROPAGATION` — anchor literal across canonical files.
- `CM-GIT-HOOKS-INSTALL-SCRIPT` — installer present + executable.
- `CM-GIT-HOOKS-SOURCE-DIR` — 4 hook bodies present + executable.
- `CM-COMMIT-ALL-SIBLING-CHECK` — `_constitution_sibling_check` in `commit_all.sh`.
- `CM-CI-WORKFLOW-PRESENT` — CI workflow + ≥ 3 required jobs.

Composes with §1.1 (paired meta-test mutations), §9 (data safety), §11.4 (end-user-quality covenant — Layer 5 CI enforcement is the §11.4 promise made mechanical), §11.4.1 (FAIL-bluffs forbidden), §11.4.4 (test-interrupt-on-discovery), §11.4.5 (audio + video quality), §11.4.6 (no-guessing), §11.4.41 (pre-force-push merge-first — this very anchor’s renumber from §11.4.74 → §11.4.75 was driven by §11.4.41 merge-first discipline), §11.4.43 (TDD-fix), §11.4.51 (live-ADB-first), §11.4.52 (autonomous-validation), §11.4.58 (PWU pipeline), §11.4.65 (universal Markdown export — Layer 1+3+4 are §11.4.65’s mechanical seam), §11.4.66 (interactive-clarification), §11.4.67 (target-shell- parseability — hooks parse under bash AND mksh), §11.4.69 (universal sink-side positive-evidence), §11.4.70 (subagent-driven execution), §11.4.71 (pre-push fetch + integrate — Layer 3 + 5 honour the merge-first pipeline), §11.4.72 (audio top-priority — hooks do not race against in-flight audio commits because they hold no lock themselves), §11.4.73 (main-spec versioning — when spec exists, hooks include it), §11.4.74 (submodule-catalogue-first — hooks can be added to the canonical catalogue for cross-project reuse).

No escape hatch. No `--skip-hooks`, `--bypass-enforcement`, `--allow-orphan-md`, `--ci-not-applicable`, `--mechanical-enforcement-not-needed` flag exists. The `--no-verify` route IS the deliberate audit-trail bypass; §11.4.75 makes the audit trail mechanical via the `commit-msg` footer requirement.

The mandate exists because the User mandate of 2026-05-20 is unambiguous: violations MUST NEVER EVER AGAIN happen. Reliance on vigilance has demonstrably failed three times in 24 hours. Only mechanical enforcement at every layer can deliver the discipline the project requires.

Propagation gate CM-COVENANT-114-75-PROPAGATION enforces this anchor literal across the ~44-file consumer fleet. Paired mutation strips the literal → gate FAILs.

Canonical authority: constitution submodule Constitution.md §11.4.75.

Non-compliance is a release blocker regardless of context.

§11.4.76 — Containers-submodule mandate (User mandate, 2026-05-20)

Forensic anchor — direct user mandate (verbatim, 2026-05-20):

“For any work or requirements of running services or codebase inside the Containers (Docker / Podman / Qemu / Emulators, and so on) we MUST USE / INCORPORATE the Containers Submodule properly: <https://github.com/vasic-digital/containers> (`git@github.com:vasic-digital/containers.git`). Containers Submodule contains all means for us to Containerize our code and services! If any feature or Containing System is missing or not supported we MUST EXTEND IT properly like we do all of our projects! All these details MUST BE added to our root constitution (constitution Submodule) - Constitution.md, AGENTS.md, CLAUDE.md, QWEN.md - and followed and respected! No bluff work is allowed of any kind!”

§-slot history note. This anchor was originally drafted as §11.4.75 but renumbered to §11.4.76 per the §11.4.71 fetch-before-push mandate after a concurrent upstream commit (0a70083) landed §11.4.75 (Mechanical Enforcement Without Exception). Same collision-resolution pattern as §11.4.75’s own renumber-from-§11.4.74.

The mandate. For ANY containerized workload — Docker, Podman, Qemu, Kubernetes, container-backed emulators (Android emulator, VM-based testbeds, etc.) — every consuming project MUST:

1. **Use the canonical Containers Submodule.** <https://github.com/vasic-digital/containers> (`digital.vasic.containers`) is the authoritative library for runtime auto-detection, endpoint discovery, lifecycle/health management, compose orchestration, cross-build, emulator integration, and on-demand service boot. Mirrored across the four §2.1-canonical hosts (GitHub + GitLab + GitFlic + GitVerse).
2. **Install as a Git submodule.** Container-orchestration concerns enter the consuming project as a git submodule `add git@github.com:vasic-digital/containers.git <path>/containers`. Project `go.mod` (or equivalent dependency manifest) consumes via `replace digital.vasic.containers => ./<path>/containers` during development; production builds resolve through pinned commit SHAs.
3. **Boot infrastructure on demand.** Tests, CLI doctor commands, and local-development workflows that depend on Postgres / Redis / OTel / message-brokers / emulators MUST invoke the Containers Submodule’s `pkg/boot` + `pkg/compose` + `pkg/health` APIs to bring infra up automatically. Operators

MUST NOT be required to manually start podman machine / docker compose up before tests — the boot is part of the test entry point. This is the **on-demand-infra invariant**.

4. **Extend the Containers Submodule, never reimplement.** Per §11.4.74's extend-don't-reimplement rule applied specifically to containerization: if a runtime (e.g., a new emulator type) or lifecycle primitive (e.g., a new health-check kind) is missing, the consuming project MUST add it to vasic-digital/containers (PR upstream + bump the consuming project's submodule pointer) — never as a parallel implementation inside the consuming project.
5. **Anti-bluff captured-evidence requirement.** Every integration test that claims to exercise a containerized component MUST actually boot that component via the Containers Submodule. Tests that pass without containers actually running (e.g., short-circuit fakes that bypass the boot path) are a §11.4 bluff violation. A passing test MUST imply the infra was up.
6. **Composition with §11.4.74.** The Containers Submodule itself is one of the Submodules surveyed by Catalogue-Check. A project tracker row touching containerization MUST record Catalogue-Check: extend vasic-digital/containers@<sha> (or reuse) — no-match is only valid if the consuming project demonstrates the Containers Submodule's API cannot model the workload (rare).

Anti-bluff captured-evidence gate (planned). CM-CONTAINERS-USED: - For every PR that touches Dockerfile*, *compose*.yaml, *podman*.yaml, Vagrantfile, or any code under a directory matching (test|integration|e2e)/.*infra heuristic, the gate scans for an import of digital.vasic.containers/.... Absence in a non-trivial container-touching PR is a FAIL. - The gate's paired mutation strips the import and asserts the gate FAILs.

Composition. Composes with §11.4.74 (catalogue-first discovery), §11.4.75 (mechanical enforcement — this clause IS one of the things mechanical enforcement protects), §3 (submodule propagation), §11.4.31 (Submodule-Dependency-Manifest), §11.4.36 (install_upstreams on clone), §11.4.28 (Submodules-As-Equal-Codebase: the Containers Submodule itself follows all dev-cycle rules), §11.4.65 (multi-format export for docs maintained inside the Containers Submodule), §1.1 (paired-mutation gates protect the Containers Submodule's own boot / health primitives).

Why this matters. Without the on-demand-infra invariant, integration tests silently degrade to “pass without infra running” (a §11.4 bluff) or break with cryptic “connection refused” errors that operators waste hours debugging. The Containers Submodule turns infra-boot into a typed, programmable, anti-bluff-tested capability rather than a sequence of imperative shell commands. Consuming projects converge on a single shared lifecycle/health model instead of each project hand-rolling their own.

Classification: universal (per §11.4.17). Applies to every project that runs code inside any containerization or virtualization runtime.

Canonical authority: constitution submodule [Constitution.md](#) §11.4.76.

Non-compliance is a process violation; landing containerized infra without consuming the Containers Submodule (i.e., reinventing compose orchestration in-project) is a release blocker.

§11.4.77 — Regeneration-mechanism-required mandate (User mandate, 2026-05-20)

Forensic anchor — direct user mandate (verbatim, 2026-05-20):

“We must be sure that after excluding anything from Git versioning we still have the mechanism which will out of the box obtain or re-generate missing content! Add this mandatory safety rule / constraint into ours root (constitution Submodule) Constitution.md, CLAUDE.md, AGENTS.md, QWEN.md or any other required document / file from the constitution Submodule! Do not forget to fetch and pull Submodule first! Commit and push all Submodule changes to all upstreams!”

Forensic incident (2026-05-20T15:00Z, parent project ATMOSphere-Android-15). The parent’s audio Tier 1 `commit_all.sh` stalled four hours on `git add -A` while scanning 274 GiB of untracked `.git-backup-20260520T084610Z-pre-orphan-kill/` + 159 GiB `RKTools/linux/` + 167 GiB `qa-results/` content — all of which should never have been tracked. The natural remediation is to add those paths to `.gitignore`, BUT doing so without a regeneration mechanism would silently orphan fresh clones: a new operator running `git clone && cd <project>` would find missing `RKTools`, missing test infrastructure, missing build outputs they cannot reproduce. §11.4.77 codifies the universal rule that closes this gap.

The mandate. Every `.gitignore` entry that excludes either (a) more than ~100 MiB of content, OR (b) any artifact essential to building / running / testing the project, **MUST** be accompanied by a documented + automated mechanism that re-obtains or re-generates the excluded content on a fresh clone. The closed-set choices are:

1. **Re-obtain** — download from an authoritative source (vendor tarball, SDK installer, npm / pip / cargo / go-mod registry, container image registry, dedicated git submodule, S3, GCS, vendor-internal artifact server) via a script that runs deterministically on a fresh clone.
2. **Re-generate** — produce the content from tracked source code via a script (build pipeline, code generation, asset rendering, captured-evidence replay, container build, kernel build, image packaging) that runs deterministically on a fresh clone.

Required artefacts per qualifying `.gitignore` entry (each is independently mandatory):

1. **Metadata file at `.gitignore-meta/<entry-slug>.yaml`** (or equivalent project-canonical location declared in the project’s `CLAUDE.md`) carrying: the `gitignore-pattern` verbatim, `mechanism-type: re-obtain | re-generate`, `script-path: <repo-relative path>` (the deterministic regeneration endpoint), `expected-disk-usage: <bytes-or-human-readable>`, `vendor-url-or-source: <URL or in-tree path>` when applicable, `integrity: { algorithm: sha256 | md5, value: <hex> }` when the content is fetchable from a known-stable mirror, `requires-network: true | false`, `requires-credentials: true | false` (and which credentials slot per §11.4.10).
2. **Entry in the post-clone bootstrap script** (`scripts/setup.sh` or the project-canonical equivalent). The script **MUST** run the regeneration mechanism non-interactively on a fresh clone or be invocable as an idempotent step (`scripts/setup.sh --regenerate <entry-slug>`).

3. **Pre-build gate** that verifies either the regenerated content is present at the expected path OR the regeneration script ran successfully (e.g., a stamp file under `.gitignore-meta/.regenerated/<entry-slug>.ok` with `mtime ≥ source freshness window`).
4. **Documentation update** — README plus the relevant docs/guides/*.md guide describing the mechanism, including the manual fallback (vendor URL, alternative download mirrors), the time + disk-usage budget, and per-§11.4.10 credentials requirements.

No escape hatch. Adding a bare `.gitignore` entry that excludes significant or essential content without the accompanying mechanism is itself a §11.4 PASS-bluff variant — the codebase appears complete to the casual eye, every pre-build / post-build gate goes green, but a fresh clone cannot build / run. There is no `--skip-regen-mechanism`, `--gitignore-is-enough`, `--operator-already-has-content` flag.

Composition with sister anchors:

- **§11.4.6 (no-guessing)** — the mechanism MUST be verified working end-to-end on a sandbox clone before the `.gitignore` entry lands. LIKELY works on a fresh clone is the exact bluff §11.4.6 forbids.
- **§11.4.65 (universal Markdown export)** — generated HTML / PDF siblings are an instance of re-generate; the auto-regeneration helper (`sync_all_markdown_exports.sh`) is the §11.4.77 mechanism for that gitignored content where applicable.
- **§11.4.66 (interactive clarification)** — if the agent is uncertain whether a candidate `.gitignore` addition qualifies under §11.4.77, the agent MUST ASK the operator via the platform's interactive question mechanism rather than guess.
- **§11.4.71 (pre-push fetch + integrate)** — the regeneration mechanism's integrity hash and vendor URL MUST be re-validated against upstream before pushing the §11.4.77-governing commit; stale hashes are a §11.4.6 violation.
- **§11.4.74 (catalogue-first + extend-don't-reimplement)** — if the regeneration mechanism can be implemented by extending a Submodule already in `vasic-digital/HelixDevelopment` (e.g., a reusable downloader / decompressor helper), the project MUST extend rather than reimplement.
- **§11.4.75 (Mechanical Enforcement Without Exception)** — the local pre-commit hook (Layer 1) MUST refuse a commit that adds new `.gitignore` lines without a corresponding `.gitignore-meta/*.yaml`; the CI workflow (Layer 5) MUST replay the gate on every push. Paired meta-test mutation: strip a metadata YAML entry → gate FAILs.
- **§11.4.76 (Containers-submodule mandate)** — container images excluded from VCS are regenerated via `vasic-digital/containers pkg/boot + Dockerfiles` tracked in-tree; the `.gitignore-meta/<entry-slug>.yaml` references the Containers Submodule build path.
- **§9 / §9.2 (zero-risk data safety)** — the regeneration mechanism MUST be pre-tested in a sandbox clone with a hardlinked-backup safety net BEFORE relying on it for live operator clones; a regeneration script that silently corrupts content on a fresh clone is a §9 violation.
- **§3 (propagation order)** — submodules consuming this constitution submodule inherit §11.4.77; their own `.gitignore` entries are subject to the same mechanism requirement, propagated via their own `CLAUDE.md / AGENTS.md / QWEN.md`.

Anti-bluff captured-evidence gate (planned). CM-GITIGNORE-REGEN-MECHANISM:

- For every PR that adds a line to `.gitignore` (or to any `*.gitignore` file in a sub-tree), the gate scans for a matching entry in `.gitignore-meta/`, verifies the `script-path` exists and is executable, verifies the metadata YAML parses and carries the required keys, and verifies the post-clone bootstrap references the mechanism.
- Bare-line additions (e.g., `echo 'qa-results/' >> .gitignore`) without the metadata sibling are a FAIL.
- The gate's paired §1.1 mutation strips one required YAML key (e.g., `script-path:`) and asserts the gate FAILs.

Why this matters. Without §11.4.77, the natural `.gitignore` ergonomics — “this is too big to track, exclude it” — leads to fresh-clone orphanage: operators or CI runners (or new developers, or end-user mirror sites) clone the project and discover essential content is missing, with no documented way to obtain it. The cost is paid at every fresh-clone event, every CI bootstrap, every operator onboarding, every disaster-recovery exercise. §11.4.77 forces the cost to be paid once, at the `.gitignore-addition` commit, where it is bounded and visible.

Forensic anchor crosswalk to today's incident. Without §11.4.77, the natural fix to today's `commit_all.sh` 4-hour-stall — adding `.git-backup-*`, `RKTools/linux/`, `qa-results/` to `.gitignore` — would have orphaned every fresh ATMOSphere clone. With §11.4.77, each of those entries now requires: (a) a `.gitignore-meta/<slug>.yaml` declaring the mechanism (e.g., `.git-backup-*` is a transient artefact safe to re-create per §9.2 `hardlinked-backup`; `RKTools/linux/` is re-obtainable from Rockchip's vendor tarball; `qa-results/` is re-generatable by re-running the on-device test cycle); (b) bootstrap-script entries; (c) gates; (d) docs. The bare `.gitignore` addition is now mechanically blocked.

Classification: universal (per §11.4.17). Applies to every project that excludes content from VCS for size / regenerability / vendor-provenance reasons.

Canonical authority: constitution submodule [Constitution.md](#) §11.4.77.

Non-compliance is a release blocker regardless of context.

§11.4.78 — CodeGraph code-intelligence mandate (User mandate, 2026-05-20)

Forensic anchor — direct user mandate (verbatim, 2026-05-20):

“All gathered knowledge and steps MUST BE added as future reusable experience into our constitution Submodule on setting up / configuring and using codegraph with all supported CLI agents we may use for our codebase! Make codegraph MANDATORY CHOICE for this purpose for all of our project and add proper comprehensive details about everything codegraph related into our root constitution (constitution Submodule) `Constitution.md`, `AGENTS.md`, `CLAUDE.md`, `QWEN.md` and other constitution Submodule relevant files! All project which do not have configured and installed codegraph yet MUST DO IT and MUST USE IT!”

The mandate. Every consuming project under this Constitution **MUST** install, initialize, and use **CodeGraph** (<https://github.com/colbymchenry/codegraph>, npm package @colbymchenry/codegraph) as its code-intelligence layer for AI coding agents. CodeGraph builds a local SQLite semantic knowledge graph of a codebase and exposes it to agents over the Model Context Protocol (MCP) — 100% local, no cloud, no external API (compatible with the Local-Only CI/CD posture). A project worked on by AI coding agents that has not yet installed and wired CodeGraph **MUST** do so; using it is not optional.

1. **Install.** `npm install -g @colbymchenry/codegraph` (Node.js 18+). The npm global prefix **MUST** be user-writable — installation **MUST NOT** require `sudo / su` (§9 / §12 host-safety + the no-privilege-escalation posture). Native modules (better-sqlite3, tree-sitter) compile on install with a WASM fallback.
2. **Initialize + index.** `codegraph init` at the project root creates `.codegraph/config.json` (tracked) and `.codegraph/codegraph.db` (a regenerable build artifact — gitignored per §11.4.30 / §11.4.77, with `codegraph index` as its declared regeneration mechanism). The `config.json` exclude list **MUST** exclude (a) other-owned submodules and vendored trees — index the consuming project's own domain code — and (b), non-negotiably, every credential/secret path (`.env*`, keystores, signing keys, service-account JSON) per §11.4.10. A secret reaching the index is a §11.4.10 violation. `codegraph index` builds the graph; `codegraph sync` keeps it fresh.
3. **Wire every supported CLI agent.** The CodeGraph MCP server (`codegraph serve --mcp`, stdio transport) **MUST** be registered with every AI coding CLI agent the project's developers use. Registration is project-scoped and committed where the agent supports it (e.g. Claude Code `.mcp.json`, OpenCode `opencode.json`, Qwen Code `.qwen/settings.json`, Crush `.crush.json`); host-local where the agent stores MCP config outside the repository (e.g. Kimi CLI `~/.kimi/mcp.json`). Every MCP config **MUST** reference the bare `codegraph` command resolved on `PATH` — never a hardcoded host path (the no-hardcoding posture). Claude Code is the canonical primary agent; the others **MUST** work too.
4. **Anti-bluff verification is mandatory.** CodeGraph integration **MUST** be covered by an anti-bluff verification suite (the canonical reference implementation is Lava's `scripts/verify-codegraph.sh + tests/codegraph/`, six layers): index reality, query correctness, MCP-protocol JSON-RPC, per-agent connectivity, per-agent end-to-end LLM drive, and a falsifiability rehearsal that mechanically proves the suite FAILs when the index is deliberately broken. The per-agent end-to-end layer **MUST** use an **unforgeable challenge** — a fact obtainable only by calling a CodeGraph MCP tool (e.g. the index node count via `codegraph_status`) — so an agent answering from its own file-reading tools cannot produce a false PASS. An agent that genuinely cannot be driven end-to-end in the test environment (missing credentials, exhausted quota, environment incompatibility) is recorded as a documented SKIP gap per §11.4.3 — never a faked PASS.
5. **Comprehensive documentation.** Every project **MUST** carry a `docs/CODEGRAPH.md` (or canonical equivalent) describing install, initialization, indexing, per-agent wiring, the verification suite, and troubleshooting — kept in sync per §11.4.12 / §11.4.65.

6. **Catalogue-first.** CodeGraph is a third-party developer tool, not an owned submodule — it does NOT enter the project as a `git submodule` (per §11.4.74 it is consumed as the published npm package) and it does NOT add a Git remote.

Anti-bluff captured-evidence gate (planned). CM-CODEGRAPH-WIRED: for every consuming project, the gate verifies `.codegraph/config.json` exists with the §11.4.10 secret-exclusions present, every developer-used agent carries a CodeGraph MCP registration, and the verification suite exists and is executable. The paired §1.1 mutation removes a secret-exclusion from `config.json` and asserts the gate FAILs.

Composition. Composes with §11.4.3 (per-environment-topology SKIP for un-runnable agents), §11.4.10 (credentials never indexed), §11.4.12 + §11.4.65 (CODEGRAPH.md kept in sync + exported), §11.4.30 / §11.4.77 (the `.codegraph/codegraph.db` artifact is gitignored with `codegraph index` as its regeneration mechanism), §11.4.74 (catalogue-first — CodeGraph is consumed, not reimplemented), §11.4 (the verification suite is anti-bluff by construction — CI green is necessary, never sufficient), §1.1 (paired-mutation gate).

Why this matters. AI coding agents that explore a codebase by repeated file scanning consume tokens and tool calls wastefully and build a shallow, drift-prone mental model. A pre-indexed semantic graph gives every agent instant, consistent symbol / caller / callee / impact resolution. Standardising on one tool — CodeGraph — across every Helix project means the capability is wired once, verified anti-bluff once, and reused everywhere, instead of each project hand-rolling ad-hoc context tooling.

Classification: universal (per §11.4.17). Applies to every project that is worked on by AI coding CLI agents.

Canonical authority: constitution submodule [Constitution.md](#) §11.4.78.

Non-compliance is a process violation; a project worked on by AI agents without CodeGraph installed, wired, and anti-bluff-verified is in breach of this mandate.

§11.4.79 — Own-org submodules MUST be included in the CodeGraph index (User mandate, 2026-05-21)

Forensic anchor — direct user mandate (verbatim, 2026-05-21):

“All Submodules we use in the project and that are part of organizations to which we have the full access via GitHub, GitLab and other CLIs MUST BE included into the codegraph database and initialized / scanned / synced!”

The mandate. This extends §11.4.78 step 2’s exclude-list refinement with a per-submodule-ownership split. Every consuming project’s `.codegraph/config.json` MUST distinguish:

Submodule class	Treatment in CodeGraph index
Own-org — full write access via the project’s CLIs (canonical orgs: <code>vasic-digital</code> on GitHub/GitLab/GitFlic/	MUST be INCLUDED. Per-submodule sources MUST be indexed so AI coding agents resolve symbols, callers, and impact across the whole

Submodule class	Treatment in CodeGraph index
GitVerse + HelixDevelopment on GitHub)	repo graph — not just the consuming project’s domain code.
Third-party — no write access, vendored only (the §11.4.74 no-match → vendor path, e.g. an external SDK like gopkg.in/telebot.v3)	MUST be EXCLUDED. Indexing third-party code wastes the local index budget and creates symbol-noise the project’s contributors cannot fix.

This refines, NOT contradicts, §11.4.78 step 2 (which spoke of “other-owned submodules” generically). The classifier is **write-access via the project’s own CLIs**, not “internal vs external” subjectively — own-org submodules are the project’s own code under another path; third-party submodules are upstream code Herald extends without owning.

Operational steps for every consuming project:

1. **Fetch + pull latest of every submodule** before re-indexing —
git submodule update --remote --merge from the project root, then verify the project still builds. Third-party submodule pins (e.g. gopkg.in/telebot.v3 pinned at v3.3.8 to keep the Go import path stable) **MUST** be respected and rolled back if a --remote advances past the load-bearing pin.
2. **Adjust .codegraph/config.json exclude to keep own-org submodules in scope.** Third-party submodule paths **MUST** be explicitly listed in exclude (per §11.4.78 step 2’s per-credential exclusion principle); own-org submodule paths **MUST NOT** appear in exclude.
3. **Re-index.** Run the project’s canonical CodeGraph initialiser (e.g. scripts/codegraph_setup.sh) so codegraph index (or codegraph sync) rebuilds the index with the corrected exclude list.
4. **Verify symbols across own-org submodules.** The project’s scripts/codegraph_validate.sh (per §11.4.78 anti-bluff verification) **MUST** probe at least one symbol that lives **ONLY** inside an own-org submodule (e.g. bg.TaskQueue from submodules/background/interfaces.go). A successful query proves the include path actually reached the submodule’s sources — not a §107 PASS-bluff where validate runs against a stale index.
5. **Paired §1.1 mutation:** temporarily add the own-org submodule path to exclude, re-index, run validate — it **MUST FAIL** on the cross-submodule probe. Restore. This mutation proves the validate is not bluffing about the include reach.

Composition. Composes with: - §11.4.74 (catalogue-first): the own-org classifier matches the reuse/extend catalogue-check disposition; third-party matches no-match → vendor. - §11.4.78 (CodeGraph mandate): this §11.4.79 refines step 2’s exclude-list contract without weakening any other step. - §11.4.10 (credentials never tracked): the exclude list **MUST** still cover every .env*, keystore, signing key, service-account JSON regardless of submodule ownership. - §1.1 (paired mutation): every validate probe added per step 5 above **MUST** land with a mutation pair. - §107 (end-user usability): an indexed graph that lies about which symbols are reachable is a §107 PASS-bluff against the AI coding agents that consume it.

Why this matters. When AI coding agents work across a project that vendors own-org capability modules (e.g. `commons_constitution` consuming the Helix-stack 9-module surface under `submodules/`), excluding those modules from the index forces the agents back into shallow file-scan exploration of half the codebase — exactly the failure mode §11.4.78 exists to prevent. The agents lose call-graph resolution across the seam between domain code and own-org infra code. Including own-org submodules is the difference between agents seeing the project as a unified graph and agents tripping over an invisible boundary at the submodule directory.

Classification: universal (per §11.4.17). Applies to every project under this Constitution that has CodeGraph configured and own-org submodules vendored.

Canonical authority: constitution submodule `Constitution.md` §11.4.79.

Non-compliance is a process violation; an AI-agent-worked project whose own-org submodules are excluded from CodeGraph (when they could be included) is in breach of this mandate. Severe cases (own-org submodules silently excluded WITHOUT an audit trail in `.codegraph/config.json` comments) are release blockers.

§11.4.80 — CodeGraph regular-update + sync automation mandate (User mandate, 2026-05-21)

Forensic anchor — direct user mandate (verbatim, 2026-05-21):

“We MUST regularly check for the updates and execute codegraph npm updates so the latest version of it is always installed on the host machine! After update executes with success use the following commands (and validate and verify them using codegraph help) to sync all: [list of codegraph commands]. Make sure we have proper full automation bash scripts which will run regularly and that these are part of the constitution Submodule since they MUST BE available to all projects which do respect and follow our constitution rules and mandatory constraints! Make sure all updates, sync processes we do and important codegraph related events are all documented under docs/codegraph in Status and Status_Summary documents (another area / context to regularly update and sync) and regularly export them like all other Status docs into the PDF and HTML!”

The mandate. Every consuming project under this Constitution MUST regularly check for CodeGraph npm-package updates, install the latest stable version, and re-sync its local index. The automation lives in the constitution submodule itself so every consuming project gets it via the inherited tree without reimplementation. Three deliverables:

1. **<constitution>/scripts/codegraph_update.sh** — checks the installed `@colbymchenry/codegraph` version against npm’s latest, runs `npm update -g @colbymchenry/codegraph` (or `npm install -g @colbymchenry/codegraph@latest`) if a newer version exists, captures the old + new version + timestamp into the constitution’s docs/codegraph/Status.md audit trail, then exits 0 only when `codegraph --version` produces the expected post-update output. Idempotent: re-runs on the same day are no-ops. Anti-bluff (§107): the “update succeeded” PASS MUST observe the new version, not just `npm update` exit code.

2. **<constitution>/scripts/codegraph_sync.sh** — after a successful update, runs the canonical sync sequence inside the consuming project (passed as an argument or auto-resolved via `find_constitution.sh` parent walk):
codegraph status (baseline) → codegraph sync . (incremental update) → codegraph status (post-sync) → `scripts/codegraph_validate.sh` (the project's anti-bluff verifier, per §11.4.78 step 4). Each step's output is appended to the project's `docs/codegraph/Status.md` ledger. The sync **MUST** consult `codegraph help` BEFORE invoking any command to verify the subcommand still exists with the same shape (tools evolve). The canonical sync sequence covered by `codegraph help`:

```
codegraph                # Run interactive installer
codegraph install        # Run installer (explicit)
codegraph init [path]    # Initialize in a project (--
index to also index)
codegraph uninit [path]  # Remove CodeGraph from a
project (--force to skip prompt)
codegraph index [path]   # Full index (--force to re-
index, --quiet for less output)
codegraph sync [path]    # Incremental update
codegraph status [path]  # Show statistics
codegraph query <search> # Search symbols (--kind, --
limit, --json)
codegraph files [path]   # Show file structure (--
format, --filter, --max-depth, --json)
codegraph context <task> # Build context for AI (--
format, --max-nodes)
codegraph affected [files...] # Find test files affected by
changes
codegraph serve --mcp    # Start MCP server
```

3. **<constitution>/docs/codegraph/Status.md + docs/codegraph/Status_Summary.md** — append-only ledgers tracking every CodeGraph-related event across the consuming project: install events, npm-update events, sync runs, index regenerations, validate runs (with PASS/FAIL counts), HRDs opened/closed against CodeGraph. Both ledgers **MUST** be exported to `.html` and `.pdf` siblings on every edit (per §11.4.65 multi-format export mandate). The `Status_Summary.md` is the operator-readable derived index (per §11.4.53 fixed-summary backfill convention).

Operational cadence. The automation **MUST** run at least weekly (per §11.4.45 status-digest cadence). Consuming projects **MAY** wire it more frequently (e.g. daily via cron or per-CI-run) but never less. Cadence is not contractually fixed because environments differ (development workstation vs CI vs production-locked host); the floor is “weekly” because slower than that risks accumulated drift.

Scripts are inherited by reference. Per §3 submodule inheritance, consuming projects do NOT copy `codegraph_update.sh` / `codegraph_sync.sh` into their own `scripts/` directory. They invoke the constitution submodule's scripts directly (e.g. `bash ${CONST_DIR}/scripts/codegraph_update.sh && bash ${CONST_DIR}/scripts/codegraph_sync.sh .`). This is the single-source-of-truth pattern — when the constitution updates the scripts, every consuming project picks up the change on its next submodule update.

Composition. Composes with §11.4.78 (CodeGraph parent mandate), §11.4.79 (own-org submodule inclusion), §11.4.10 (credential exclusions stay applied), §11.4.45 (status-digest cadence), §11.4.53 (Fixed_Summary backfill), §11.4.65 (multi-format export), §107 (each sync run’s “ok” MUST carry observed-version evidence, not just exit code), §1.1 (paired mutation: temporarily downgrade installed codegraph version → script MUST detect drift and self-correct → restore).

Why this matters. A CodeGraph index built against an older version of the indexer may miss symbols that newer versions resolve correctly. Stale indexes silently lie to AI agents about what symbols exist and how they connect. Regular automated updates + sync close the loop: tooling stays current, the index reflects current source, AI agents see the truth. Manual “run codegraph index when I remember” is exactly the failure mode this mandate forbids.

Classification: universal (per §11.4.17). Applies to every project under this Constitution that has CodeGraph configured per §11.4.78.

Canonical authority: constitution submodule Constitution.md §11.4.80.

Non-compliance is a process violation; severe cases (consuming project has not run `codegraph_update.sh` in >2 weeks AND has open AI-agent work) are release blockers — the agents’ index reflects a stale tool version and may produce silently wrong reasoning.

§11.4.81 — Cross-platform-parity mandate (User mandate, 2026-05-21)

Forensic anchor — direct user mandate (verbatim, 2026-05-21):

“Any Linux-only blocker / issue we have MUST BE created macOS and other supported platforms equivalent! So, depending on platform proper implementation will be used for particular OS! EVERYTHING MUST BE PROPERLY EXTENDED AND UPDATED!”

The mandate. Every consuming project whose supported-platforms manifest lists more than one OS MUST, for every feature/test/gate/challenge/mutation that depends on platform-specific primitives, ship a per-OS-equivalent implementation chosen at runtime via `uname -s` (or equivalent platform detection — `[[ $OS == 'Windows_NT' ]], sw_vers -productName, /etc/os-release` on Linux distros, etc). A Linux-only gate that ships without macOS / Windows / BSD equivalents (when the kernel of those OSes permits one) is a §11.4 / §107 PASS-bluff at the platform-coverage layer: the test reports GREEN on Linux but tells nothing about whether the feature works for end users on every supported platform.

Three sub-mandates.

(A) Per-OS implementation REQUIRED. Every feature whose Linux implementation uses a `cgroup` / `systemd` / `proc` primitive MUST have a documented per-OS equivalent (POSIX `setrlimit` / `ulimit`, macOS `launchd plist` `HardResourceLimits`, BSD `rctl`, Windows Job Object, etc) chosen via runtime platform detection. The wrapper/library code MUST dispatch to the correct branch with NO operator awareness required — the same operator-path invocation works on every supported platform. Canonical example from `vasic-digital/tmux`: `tmx`

new -s NAME dispatches to systemd-run --user --scope on Linux and to scripts/tmx-rlimit-wrapper.sh (POSIX setrlimit) on Darwin — same operator command, OS-correct mechanism per session.

(B) Per-OS tests REQUIRED. Every gate test that exercises platform-dependent behavior MUST have a per-OS branch under case "\$(uname -s)" in. Each branch MUST exercise the platform's actual native primitive with positive captured evidence per §11.4.2 + §11.4.5 (e.g. read ulimit -t inside the session via send-keys + capture-pane on Darwin to verify RLIMIT_CPU applied; read /sys/fs/cgroup/.../memory.max content on Linux to verify cgroup MemoryMax applied). Platform-conditional SKIP-with-reason is acceptable ONLY when the platform genuinely cannot enforce the invariant at all (kernel limit — see (C) below), AND the SKIP message MUST cite the precise kernel/system limitation with a reproducer + a link to the project's honest-gap section in its operator guide (typically docs/guide/README.md or equivalent).

(C) Honest kernel-gap citation + adjacent equivalent test REQUIRED. Where a Linux primitive has NO macOS / other-OS equivalent due to a documented kernel/OS limitation (canonical example: XNU does NOT enforce RLIMIT_AS / RLIMIT_DATA / RLIMIT_RSS for unprivileged processes — verified by reproducer in vasic-digital/tmux docs/guide/README.md §5.6), the test MUST: (1) detect the gap at runtime, (2) SKIP with the exact kernel reason + reproducer + link to the project's honest-gap doc, (3) **provide an ADJACENT test that exercises the closest invariant the platform CAN enforce** (e.g. RLIMIT_CPU + SIGXCPU as the macOS proxy for “process is bounded under load”). The adjacent test MUST itself be anti-bluff per §11.4 with positive captured evidence and a paired §1.1 mutation.

Composes with. §11.4.1 (FAIL-bluffs equally forbidden — a Linux-PASS / Darwin-SKIP-without-honest-reason is a §11.4.81 violation), §11.4.2 (recorded-evidence — each branch's captured evidence is platform-specific), §11.4.3 (per-host-topology dispatch — §11.4.81 strictens §11.4.3 by requiring an equivalent test when the platform CAN enforce, not just SKIP-with-reason for genuine kernel gaps), §11.4.5 (captured evidence quality — platform-specific quality analysis), §11.4.4 (test-interrupt-on-discovery — discovering a Linux-only test without a Darwin equivalent triggers §11.4.81 mid-cycle), §11.4.6 (no-guessing — “platform X probably can't do Y” is a §11.4.6 violation; prove via reproducer or mark UNCONFIRMED), §11.4.20 / §11.4.70 (subagent-driven — cross-platform branches are independent enough to dispatch as parallel subagent work per branch), §11.4.27 (no-fakes-beyond-unit + 100% test-type coverage — platform branches inherit the 100% type-coverage discipline), §11.4.69 (universal sink-side positive-evidence taxonomy — every platform branch must produce a taxonomy-mapped sink-side evidence artefact), §107 (end-user usability — multi-platform projects whose tests cover only one platform deliver no end-user guarantee on the others).

Per-OS implementation equivalence catalogue (canonical examples — not exhaustive).

Linux primitive	Darwin equivalent	BSD equivalent	Windows equivalent
systemd-run --user --scope (transient cgroup)	scripts/<rlimit-wrapper>.sh using ulimit -t (RLIMIT_CPU) +	rctl process:0:deny:vmemoryuse (root); ulimit -v (unprivileged limited)	Windows Job Object via CreateJobObjectW + SetInformationJobOb

Linux primitive	Darwin equivalent	BSD equivalent	Windows equivalent
	ulimit -u (RLIMIT_NPROC)		
cgroup MemoryMax enforcement	XNU GAP — unprivileged RLIMIT_AS returns EINVAL; root-only via launchd HardResourceLimits plist. Adjacent test: RLIMIT_CPU + SIGXCPU proxy	rctl process:0:deny:memoryuse (root)	Job Object JOB_OBJECT_LIMIT_PR
cgroup TasksMax	RLIMIT_NPROC via ulimit -u (per-user, kernel-enforced)	rctl user:NN:deny:maxproc	Job Object JOB_OBJECT_LIMIT_AC
/proc/<pid>/ oom_score_adj	No equivalent — Darwin/BSD have no kernel OOM killer (different memory model). SKIP-with- reason; adjacent: rlimit bounds rogue session reach	/dev/oom-style not present in stock BSD	None at OS level (proce via Job Object)
/sys/fs/ cgroup/... introspection	ulimit -a readback inside session + ps ancestry	procstat -r <pid>	QueryInformationJob
kill -KILL on cgroup process to test scope isolation	kill -KILL on tmux server pid directly; verify sibling servers survive via direct socket query	same	TerminateProcess + p check

Operational discipline. When a consuming project adds a new feature whose Linux path lands first, the SAME pull request MUST add the per-OS branches for every other platform in the project’s supported-plat-forms manifest (or explicitly classify the platform per (C) — honest gap + adjacent test). Skipping the per-OS branches “for later” is itself a §11.4.81 violation: the gap accrues silently and the project ships Linux-PASS / other-OS-untested code while claiming multi-platform support.

Pre-build gate CM-CROSS-PLATFORM-PARITY (when implemented in the consuming project): scans every gate test for case "\$ (uname -s)" blocks, asserts a non-SKIP branch (or honest-gap citation per (C)) exists for each platform in the supported-platforms manifest. Paired §1.1 mutation: strip the Darwin branch from any one test → gate FAILs.

Classification: universal (per §11.4.17). Applies to every project under this Constitution whose supported-platforms manifest lists more than one OS. Single-platform projects are unaffected (no parity gap exists).

Canonical authority: constitution submodule [Constitution.md](#) §11.4.81.

Non-compliance is a release blocker on multi-platform projects. No escape hatch — no `--linux-only-acceptable`, `--no-platform-parity-required`, `--platform-gap-just-skip` flag exists. A Linux-PASS / other-OS-untested feature in a multi-platform project is a §11.4 / §107 PASS-bluff at the platform-coverage layer.

§11.4.82 — Iteration-speedup discipline mandate (User mandate, 2026-05-22)

Forensic anchor — direct user mandate (verbatim, 2026-05-22):

“How can we speed-up this whole development and fixing process? ... Do not forget to all speed optimizations critical rules and mandatory constraints MUST BE all added into our root (constitution Submodule) Constitution.md, CLAUDE.md, AGENTS.md and QWEN.md and all other relevant constitution Submodules files!”

Iteration cycle time is a first-order quality enabler. A 40-min rebuild cycle that catches one defect per cycle bounds the project’s defect-discovery rate at 1 defect per 40 min; a 15-min cycle bounds it at $\sim 3\times$ that rate. Slow cycles are a §11.4 PASS-bluff at the velocity layer — the project ships fewer validated features per unit of operator time than it should.

The 2026-05-22 forensic session that produced this anchor witnessed (a) 2 subagent watchdog stalls each ~ 10 min wasted, (b) 1 stalled `commit_all.sh` ~ 10 min lost on submodule cascade prompt, (c) 1 wrong-call-site speculative kernel patch (§GQ.2) burning a full 45 min rebuild cycle before forensic evidence pointed to the correct site (§GT Fix B-1), (d) 1 stale `git index.lock` from a killed process blocking the next commit, (e) 1 `auto-git gc` maintenance colliding with urgent commit. Each line item cost between 5 and 45 minutes of wall-clock that should have been productive iteration time. The aggregate impact for that single morning was ~ 3 hours of operator wait dilated onto a job whose intrinsic compute was ~ 90 minutes.

The mandate. Every consuming project’s build / test / commit / debug pipeline MUST adopt the following speedup disciplines AS MANDATORY. Each is independently enforceable:

(A) Phase 1 forensic before any speculative source patch

Before applying ANY non-trivial source patch (kernel, framework, native code, HAL, ASoC), the agent MUST first complete Phase 1 of `superpowers:systematic-debugging` — read the error/log/call-chain end-to-end and identify a FACT-grade root cause. Speculative patches without Phase 1 evidence are §11.4.6 (no-guessing) violations AND §11.4.82 violations. A wrong-call-site rebuild cycle wastes 30-90 min of compute; the alternative (15-30 min of focused source reading) is a net positive every time. Phase 1 is mandatory regardless of operator time pressure — “we don’t have time for forensics, just rebuild” is the §11.4.82 anti-pattern.

(B) Live-ADB-First (or live-equivalent) before any rebuild

Per §11.4.51 — strengthened by §11.4.82 to a release-blocker mandate. ~70% of audio / UI / boot-script / config changes are LIVE_ADB_TESTABLE (push + setprop + pm install -r + 30-second validation cycle). Only kernel / framework Java/AIDL / native C++-in-APEX / sepolicy / init.rc / ro.* / Android.bp / XML-resource-overlays / codec-XML-in-APEX changes are genuinely REQUIRES_REBUILD. Skipping the live-probe step for a LIVE_ADB_TESTABLE change is a §11.4.82 violation — costs operator 45 min of rebuild for what could have been 30 seconds of adb push.

(C) Pre-flight before launching rebuild orchestrators

Before any nohup bash scripts/<rebuild>.sh & invocation, the agent MUST run a ~30-second pre-flight that verifies (i) target device(s) reachable via ADB, (ii) downstream sinks (Arvus AVR, HDMI sink, Bluetooth peers) reachable as the test expects, (iii) host has sufficient memory and disk per §12.6 budget, (iv) no stale lock files in .git/ or build directories, (v) no orphan git/build processes from prior killed iterations. A rebuild orchestrator launched against a broken precondition wastes 45 min discovering at the test stage what 30 seconds of pre-flight would have caught.

(D) Persistent build caches outside containers

Containerized AOSP / kernel builds MUST persist ccache AND any equivalent build-system intermediate cache (Soong cache, sccache, Gradle daemon state) on the host via bind-mount, NOT inside the ephemeral container layer. Without this, every container restart drops the cache and forces a cold rebuild (~40 min for AOSP+kernel on a fast workstation). With it, incremental rebuilds where only one driver / one app changed drop to 5-15 min (ccache hit rates 80-95%). This is a one-time setup of ~15 min for ~25-32 min savings per subsequent cycle.

(E) Module-only rebuild for loadable-module-only changes

When a patch touches ONLY one or more files inside kernel/.../drivers/*/foo.c AND the affected drivers are built as loadable modules (CONFIG_*=m in the active defconfig), the rebuild MUST be make -C kernel M=<driver-path> modules (~2-5 min) NOT a full kernel rebuild (~25 min). A flash_kernel_only.sh or equivalent helper writes the rebuilt module(s) to the running device without re-flashing system.img. This saves 15-20 min per cycle for the ~10-30% of kernel patches that qualify. Built-in drivers (CONFIG_*=y) do not qualify — full rebuild required for those.

(F) Parallel multi-device testing

When the project owns more than one validation device (D3, D4, ... in topology terms), the test orchestrator MUST run autonomous validation cycles on EVERY device in parallel via background processes, with separate qa-results/<TS>/<device-tag>/ output directories per device. This catches per-device topology defects (per §11.4.3) one cycle earlier instead of waiting for a device-N-focused re-cycle. Direct wall-clock savings = 0 (parallel); amortized savings = 5-15 min per cycle when a device-specific defect surfaces.

(G) Subagent scope discipline + worktree isolation

Subagents dispatched per superpowers: subagent-driven-development MUST be: - **Scope-bounded** to ≤ 30 min of focused work with intermediate output emitted as the subagent progresses (the canonical watchdog stalls at 600s of no-output; if a subagent writes its findings file incrementally it survives even when blocked on an external dependency). - **Worktree-isolated by default** via the Agent tool's isolation: "worktree" parameter when supported, so parallel subagents do not contend on the same working tree and so a subagent crash never corrupts the conductor's state. - **Single-responsibility** — one investigation OR one fix-implementation OR one doc draft, not multiple. The empirical pattern from the §GT investigation is that 8-task subagents stall while 2-3-task subagents complete reliably.

(H) Lock-file + stale-process hygiene

Killed or crashed agents leave behind .git/index.lock, .git/.commit_all.lock, .lock files in build directories, and orphan child processes (git add -A, git pack-objects, git gc). The next agent MUST detect these on session start (or after operator-visible crash) and clean them — never wait silently for a never-arriving release. Auto git-gc MUST be disabled in repos with concurrent multi-agent work (git config gc.auto 0) — its background invocation is a guaranteed lock-contention source.

(I) Cycle telemetry per §11.4.24 build-resource stats

Every iteration cycle MUST record (i) commit hash of the source under test, (ii) wall-clock of build + flash + test phases separately, (iii) the speedup-discipline flag set (which of A-H applied), (iv) outcome (PASS / FAIL / partial / stall). Aggregated weekly, this exposes which speedup gives the biggest empirical return on the specific project, and which is being neglected. The §11.4.24 build-resource stats tracker is the canonical telemetry pipeline; §11.4.82 extends its row schema with iteration-discipline columns.

Composes with. §11.4.4 (test-interrupt-on-discovery + retest), §11.4.6 (no-guessing — speculative patches forbidden), §11.4.9 (batch-source-fixes-before-rebuild), §11.4.20 / §11.4.70 (subagent-driven), §11.4.24 (build-resource stats), §11.4.42 (iteration-discipline conductor loop), §11.4.43 (TDD-fix discipline — RED-test enables faster cycle reads), §11.4.50 (deterministic consistency — caches don't break determinism if ccache config disables time-sensitive headers), §11.4.51 (LIVE-ADB-first maximisation — strengthened to mandate here), §11.4.52 (autonomous-validation — every cycle's verdict is automation-produced), §11.4.58 (parallel-development PWU — parallel devices and parallel subagents are §11.4.58 instances), §12.7 (-j2 host-safety cap remains — speedups MUST stay within memory budget), §107 (end-user usability — fast iteration = fewer cycles → fewer defects shipped).

Pre-build gate. CM-ITERATION-SPEEDUP-DISCIPLINE (when implemented in consuming project): audits the most recent N iteration cycles for telemetry entries citing which of (A)-(I) applied; flags any cycle where none of (A)-(I) was active as a candidate for review. Paired §1.1 meta-test mutation: strip the speedup-flag column from a synthetic telemetry row → gate FAILs.

Operational discipline. When a consuming project ships an iteration speedup (e.g. lands a persistent-cache bind-mount, a parallel-device orchestrator, a module-only rebuild helper), the change **MUST** be classified as universal (per §11.4.17 — most are) OR project-specific (rare — e.g. a script that hardcodes D3/D4/Arvus IP 192.168.4.185). Universal speedups belong in this constitution submodule’s helper inventory (or in a separate Containers submodule per §11.4.76); project-specific wrappers stay in the consuming repo.

Classification: universal (per §11.4.17). Applies to every project under this Constitution. The disciplines are language-/platform-/build-system-agnostic.

Canonical authority: constitution submodule Constitution.md §11.4.82.

Non-compliance is a release blocker. No escape hatch — no `--skip-phase1-forensic`, `--no-pre-flight`, `--rebuild-everything-always`, `--unlimited-subagent-scope`, `--ignore-locks`, `--no-telemetry` flag exists. The disciplines exist because their absence has been forensically demonstrated to cost the operator hours per session.

§11.4.83 — docs/qa/ end-user evidence mandate (User mandate, 2026-05-22)

The mandate. Every feature that ships **MUST** carry a recorded end-to-end communication transcript plus any attached materials (screenshots, request/response payloads, audio, file uploads) committed under `docs/qa/<run-id>/` — one directory per feature run. A feature with no QA transcript is itself a §11.4 / §107 PASS-bluff: it claims to work but carries no auditable runtime evidence that an end user actually exercised it through the same interface they will use in production.

Forensic anchor (verbatim user mandate, 2026-05-22):

“every feature that ships **MUST** carry a recorded e2e communication transcript + any attached materials under `docs/qa/<run-id>/` (per-feature subdirectories). A feature with no QA transcript is itself a §107 PASS-bluff — it claims to work but has no auditable runtime evidence. Bot-driven automation **MUST** preserve full bidirectional communication threads as proof.”

Operative rule.

1. Every consuming project **MUST** maintain a `docs/qa/` tree (no exception, no escape hatch). Each new feature run lands under `docs/qa/<run-id>/` where `<run-id>` is monotonic + greppable (timestamp, HRD-NNN, ATM-NNN, or other workable-item identifier — see §11.4.54).
2. The transcript **MUST** be full bidirectional — every prompt/command sent + every response received + every error message + every state change observed. One-sided (“we sent X”) is not a transcript; both halves are required.
3. Attached materials **MUST** be committed alongside (screenshots in `.png`, payloads in `.json/.toon/.txt`, audio in `.wav/.mp3`, etc.). External-only links (Slack URL, Drive URL) are §11.4.13 sink-side violations — the evidence **MUST** live in-repo.
4. Bot-driven / agent-driven automation (a consuming project’s QA bot, e.g. Herald’s planned qaherald, or analogous binaries in other projects) **MUST**

preserve the full conversation thread as the proof artefact. A bot that runs the round-trip but stores only the final PASS/FAIL line is itself a §107 bluff at the QA-automation layer.

5. CI / release gates MUST refuse to tag a version that has any feature-shipping commit without its matching docs/qa/<run-id>/ directory present and non-empty. The release script in the consuming project enforces this; the constitution submodule's scripts/qa_coverage_audit.sh (when implemented) gives the universal scanner.

Composes with §11.4.2 (recorded-evidence requirement — docs/qa/ is the canonical capture location for end-user-facing features), §11.4.5 (captured-evidence quality analysis applies per-transcript), §11.4.13 (sink-side evidence stays in-repo, not behind external links), §11.4.65 (multi-format export — transcripts may be exported to PDF/HTML siblings when operator-readable), §11.4.69 (universal sink-side positive-evidence taxonomy — docs/qa/ is the “user-facing-channel” branch), §107 (an end-user feature without an end-user-channel transcript is a §107 PASS-bluff), §1.1 (paired mutation: delete a docs/qa/<run-id>/ → release gate FAILs).

Classification: universal (per §11.4.17). Every project under this Constitution carries this rule. Single-feature pre-implementation projects are exempt until their first user-visible feature lands.

Canonical authority: constitution submodule Constitution.md §11.4.83.

Non-compliance is a release blocker. No --qa-evidence-optional, --qa-transcript-later, --qa-bot-summary-suffices flag exists.

§11.4.84 — Working-tree quiescence rule for subagent commits (User mandate, 2026-05-22)

Short tag: working-tree quiescence.

The mandate. No subagent commit may proceed while any concurrent mutation gate, paired-mutation experiment, or other in-flight mutation is live in the same checkout. Before `git add`, the committing agent MUST `grep` its own working tree for mutation markers (MUTATED for paired, // always pass, return `json.Marshal` shortcut paths, // MUTATION / # MUTATION annotations, `_mutated_*` filename suffixes, etc.) and explicitly account for every modified file in the staging area. Any unexplained file in the staging area triggers ABORT.

Forensic anchor (verbatim user mandate, 2026-05-22):

“no subagent commit may proceed while any concurrent mutation gate is in flight in the same checkout. Before `git add`, the committing agent MUST `grep` its own working tree for mutation markers (MUTATED for paired, // always pass, return `json.Marshal` shortcut paths, etc.). Any unexplained file in the staging area triggers ABORT.”

Lesson (forensic case study). A consuming project's logo-fix subagent (Herald commit 72e81ab, 2026-05-21) ran in a checkout where a paired §1.1 mutation gate had temporarily introduced an // always pass shortcut into `commons_auth/middleware.go` (JWT-bypass mutation, intended to be reverted by the same gate). The subagent's `git add + git commit` swept the mutation residue into the same

commit as the unrelated logo fix, and the resulting commit was pushed to all four mirrors before any other agent caught it. The fix (Herald d5bd360, “SECURITY FIX: restore commons_auth/middleware.go JWT verify”) landed within the hour, but the window during which production-equivalent binaries shipped with a bypassed JWT verify is a real security-defect window — small but non-zero — and is the canonical example of why this rule exists.

Operative rule.

1. Subagent (and main-thread) commit flows **MUST** run a `pre-git add` quiescence check that:
 - Greps the working tree for mutation markers (canonical list above; the consuming project **MAY** extend with project-specific markers).
 - Lists every modified / untracked / staged file and matches each one against the subagent’s declared scope. Any file outside the declared scope → **ABORT**.
 - Cross-checks `git status --porcelain` against the subagent’s task description; unaccounted entries → **ABORT**.
2. Any active mutation gate (paired-mutation experiment in progress, AVR/AVI mutation cycle, manual `// always pass patch`, etc.) **MUST** be serialised — the gate runs to completion (mutate → assert **FAIL** → restore → assert **PASS**) and the working tree is verifiably clean **BEFORE** any unrelated commit may proceed in the same checkout.
3. Concurrent subagents working in the **SAME** checkout (single working tree, no `git worktree add`) **MUST** coordinate through a lockfile or marker file (e.g. `.git/MUTATION_IN_PROGRESS`) so a logo-fix subagent cannot race a JWT-mutation subagent. The constitution submodule’s `scripts/quiescence_check.sh` (when implemented) is the universal lock helper.
4. When parallel work is necessary, the §11.4.20 / §11.4.70 subagent-driven mandate **SHOULD** be combined with `git worktree add` to give each subagent an isolated working tree — eliminates the cross-mutation race by construction.
5. Post-commit, an automated mutation-residue-scanner (consuming project’s `scripts/mutation_residue_audit.sh`) **MUST** run before push. Any commit containing a mutation marker → push **BLOCKED**, commit **MUST** be reverted or amended before mirrors are updated.

Composes with §1.1 (mutation-paired gates — quiescence rule protects the mutation cycle from concurrent contamination), §11.4.20 / §11.4.70 (subagent-driven-by-default — quiescence rule is what makes parallel subagent dispatch safe), §11.4.27 (no-fakes-beyond-unit — a mutation residue swept into a commit **IS** a fake-pass surface in production), §11.4.10 (credentials handling — same class of “do not let unrelated content leak into a commit”), §107 (a security-bypass mutation that ships to production is the gravest §107 **PASS**-bluff: tests pass against the mutated codebase, but the end user loses an entire security property), §11.4.71 (pre-push fetch + investigate — quiescence check is the pre-push companion).

Classification: universal (per §11.4.17). Every project under this Constitution carries this rule.

Canonical authority: constitution submodule [Constitution.md](#) §11.4.84.

Non-compliance is a release blocker. No `--allow-residue`, `--skip-quiescence`, `--mutation-cleanup-later` flag exists. A mutation marker that lands in a tagged commit is a critical defect regardless of how briefly it persisted.

§11.4.85 — Stress + Chaos Test Mandate (User mandate, 2026-05-24)

Short tag: stress-chaos-mandate.

Forensic anchor (verbatim user mandate, 2026-05-24):

“Every fix or improvement you do MUST BE covered with full automation stress and chaos tests so we are sure nothing can break the functionality and all edge cases are monitored and polished and additionally fixed if that is needed! Everything must produce rock solid proofs and follow fully no-bluff policy!”

The mandate. Every fix or improvement landed in a consuming project MUST ship with full-automation **stress** AND **chaos** test suites that exercise edge cases, sustained load, concurrent contention, and failure-injection. A fix that PASSES its happy-path test but has never been exercised under stress or under fault-injection is a §11.4 / §107 PASS-bluff at the resilience layer: it claims to work but carries no evidence that real-world adversarial conditions (sustained throughput, parallel contention, partial failure, resource exhaustion, malformed input) leave the fix intact and the user-visible behaviour correct.

Definitions (closed-set, mechanically auditable):

1. **Stress test** — exercises the fix-under-test under sustained or concurrent load above ordinary usage. At MINIMUM one of:
 - **Sustained load** — $N \geq 100$ sequential iterations OR ≥ 30 seconds wall-clock continuous load. Per-iteration latency MUST be recorded; percentile distribution (p50/p95/p99) MUST be reported.
 - **Concurrent contention** — $N \geq 10$ parallel invocations. All N MUST complete. No deadlock, no resource leak (file-descriptor count, process-table count, RSS), no data race in shared state.
 - **Boundary conditions** — minimum input (empty), maximum input (provider/codec/protocol-limit-bound), boundary input (off-by-one at every size threshold). Each boundary MUST produce a categorised result (success, categorised-error, deterministic-skip) — NEVER an uncaught exception or silent corruption.
2. **Chaos test** — exercises the fix-under-test under failure-injection per the §11.4.69 closed-set evidence taxonomy. At MINIMUM one chaos category appropriate to the fix-class:
 - **Process-death injection** — kill the primary process, dependency process, or upstream service mid-operation. Recovery path MUST be deterministic + categorised (clean error, automatic restart, circuit-breaker open).
 - **Network-fault injection** — drop / delay / reorder packets on the relevant interface mid-call. Categorised error per §11.4.69 (e.g. category=network / category=upstream).
 - **Input-corruption injection** — corrupt the input file / config / .env mid-test. Test MUST detect + report the corruption — NEVER silently consume corrupted input.
 - **Resource-exhaustion injection** — fill disk to 99 %, allocate memory pressure (OOM-injection), exhaust file descriptors, exhaust sockets. Fix MUST refuse new work cleanly OR degrade gracefully — NEVER crash with uncaught error.

- **State-corruption injection** — mid-flight database lock loss, mid-flight file-system partial-write error, mid-flight cache invalidation. Recovery path MUST restore consistent state.

Anti-bluff (mandatory):

1. Every stress + chaos test PASS MUST cite a captured-evidence artefact path per §11.4.5 + §11.4.69. Acceptable evidence: `per-iteration_latency.json` / `throughput.csv` / `categorised_errors.txt` / `state_delta_snapshot.json` / `process_lifecycle.log` / `recovery_trace.log`. Metadata-only PASS (“all iterations exited 0”) without per-iteration latency / failure / state evidence is itself a §11.4 PASS-bluff at the stress-layer.
2. Helper library: every consuming project SHOULD ship a `stress_chaos.sh` (or analogous language-binding library) exposing reusable primitives — `ab_stress_run`, `ab_stress_concurrent`, `ab_chaos_kill_pid_during`, `ab_chaos_drop_network_during`, `ab_chaos_corrupt_file_during`, `ab_chaos_oom_pressure_during`, `ab_chaos_disk_full_during`. Each helper composes with `ab_pass_with_evidence` / `ab_skip_with_reason` per §11.4.69.
3. Chaos-injection cleanup is non-negotiable. A test that corrupts a `.env` MUST restore it in `trap '...' EXIT`. A test that fills the disk MUST `rm` the filler in `EXIT`. A test that kills a process MUST verify the process is restarted (or explicitly leave it down with operator-visible reason). Cleanup failure = §11.4.14 (test-playback cleanup) violation.

4-layer coverage per §11.4.4(b):

- **Pre-build gate** — stress + chaos test files exist + executable + parseable under `sh -n + bash -n` per §11.4.67; helper library exists + executable; the fix’s pre-build gate cites the stress + chaos test file path.
- **Paired meta-test mutation** — per §1.1, removing the chaos-injection step or stripping the per-iteration evidence-capture → the gate FAILs.
- **On-device test** (if the fix is `LIVE_ADB_TESTABLE` per §11.4.51) — the stress + chaos test is dispatched against a real device, captured-evidence directory under `qa-results/<run-id>/stress_chaos/`.
- **HelixQA Challenge entry** (if the fix is a user-visible feature per §11.4.4(b) layer 4) — a Challenge bank entry references the same stress + chaos test as the validation route.

Composition with related §11.4 anchors:

- **§11.4 / §107** — stress + chaos test PASS is the strongest end-user-quality signal: it demonstrates the fix survives adversarial conditions the happy-path test cannot reach.
- **§11.4.1** — FAIL-bluffs forbidden. A stress test that crashes with `set -u` because the chaos-injection step left a variable unset is a FAIL-bluff; fix at source.
- **§11.4.5** — captured-evidence content quality applies to stress + chaos evidence too (latency distribution recorded, not just count; error categories enumerated, not just “some errors”).
- **§11.4.6** — no guessing. Categorised errors per closed-set. “Probably the network” is forbidden — capture the categorisation.
- **§11.4.43** — TDD RED-first. The stress + chaos suite SHOULD start RED: write the test, observe the failure under load/chaos, land the fix, observe the GREEN.

- **§11.4.50** — deterministic consistency. The stress test’s N iterations MUST produce identical exit codes AND identical evidence-hashes (volatile prefixes stripped) per §11.4.50.
- **§11.4.52** — autonomous validation. Stress + chaos tests MUST run end-to-end without operator presence.
- **§11.4.69** — universal sink-side positive-evidence taxonomy. Every stress + chaos PASS uses `ab_pass_with_evidence` with a real captured-evidence file path.
- **§11.4.83** — docs/qa/ end-user transcript discipline composes with chaos-test evidence (chaos-test recovery transcripts ARE end-user-channel proofs).

Classification: universal (per §11.4.17). Every consuming project carries this rule. Stress + chaos discipline is language-/platform-/build-system-agnostic.

Canonical authority: constitution submodule Constitution.md §11.4.85.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--skip-stress`, `--no-chaos`, `--happy-path-suffices`, `--stress-test-later` flag exists. The discipline exists because the User mandate is unambiguous: “Everything must produce rock solid proofs and follow fully no-bluff policy!” — and a fix never tested under stress + chaos is not a rock-solid proof.

§11.4.86 — Roster/corpus-backed Status-doc auto-sync mandate (User mandate, 2026-05-25)

Forensic anchor — verbatim user mandate (2026-05-25):

“Make sure that assets and players Status docs are ALWAYS regularly updated and in sync like all others Status docs — any time we add or modify the assets content(s) or we change or add new / remove existing pre-installed video and audio player apps! This MUST WORK OUT OF THE BOX!”

Some Status docs (§11.4.45) are backed not by hand-written narrative alone but by a **tracked roster** (a set of installed apps / components) or a **tracked asset corpus** (a directory of test / media assets). For these, freshness cannot be left to operator vigilance — the moment a member of the roster/corpus changes (a player app added / removed / renamed; a test asset added / modified / removed) the Status doc, its Status_Summary, and their HTML + PDF exports MUST be brought back into sync **out of the box**, mechanically, not by remembering to.

The mandate. Every Status doc whose subject is a tracked roster or asset corpus MUST be kept in sync with that subject by a mechanism where **all** of the following hold:

1. **Drift-proof signal.** A content *fingerprint* — sha256 of the sorted member list, NOT mtime (which `git checkout` resets) — is persisted in a sidecar next to the Status doc. Adding / removing / modifying a member changes the fingerprint deterministically.
2. **Sync helper.** A helper regenerates the fingerprint and re-exports HTML + PDF (composing with the §11.4.65 universal exporter). Invocable standalone AND wired so the sync happens automatically (commit path and/or pre-build).
3. **Enforcement gate.** A pre-build gate recomputes the live fingerprint and FAILs when it differs from the persisted one — forcing the Status doc to be updated

whenever the roster/corpus changes. Mirrors §11.4.12 CM-ISSUES-SUMMARY-SYNC + §11.4.45 sync_integration_status.

4. **Paired §1.1 mutation.** A meta-test mutation corrupts the fingerprint (introducing drift without a doc update) and asserts the gate FAILs — proving the gate is not itself a bluff.

Composes with §11.4.12 (auto-generated docs sync — the canonical sibling pattern), §11.4.45 (integration-status-doc — §11.4.86 is its roster/corpus specialisation), §11.4.53 + §11.4.56 (Summary parity), §11.4.57 + §11.4.59 (README doc-link freshness), §11.4.60 (composite doc sync), §11.4.65 (universal Markdown export), §11.4.6 (no-guessing — the fingerprint is FACT, not a guess about freshness).

Classification: universal (per §11.4.17) — the principle (roster/corpus-backed Status docs auto-sync on subject change) is project-agnostic; the consuming project supplies the specific docs, roster/corpus sources, helper, and gate name in its own CLAUDE.md / AGENTS.md / QWEN.md per §11.4.35.

Canonical authority: constitution submodule Constitution.md §11.4.86.

Non-compliance is a release blocker regardless of context. No escape hatch — no --skip-roster-sync, --allow-status-drift, --roster-sync-not-applicable flag exists.

§11.4.87 — Endless-loop autonomous work + zero-idle agent dispatch + anti-bluff testing mandate (User mandate, 2026-05-26)

Forensic anchor — verbatim User mandate (2026-05-26):

“when we say to you that all work MUST BE continued in endless loop until there is no any open items, no unfinished workable items from our Issues docs, or from Continuation document or any unfinished work by agents at the moment it means that you will ALWAYS run until there is nothing workable elft to be done in endless loop, fullyautonomously. You will spawn agents or agents-driven work whenever that is possible or required! Not a single agent or main work stream will sit idle except if it waits for the results of something - some execution and similar! All work MUST BE always covered with comoprehensive tests - all supported tests for the project which produce real proofs for the contexts and functionalities they are testing! Bluff of any kind is not allowed and all work and tests MUST WORK completely as proofs (phisical) driven and in complete anti-bluff manner!”

When an operator instructs an AI agent to “continue in endless loop fully autonomously” — or any semantically-equivalent phrasing — the agent MUST interpret this as a HARD-CONTRACT execution covenant covering five non-negotiable obligations:

(A) Endless-loop continuation. The agent MUST continue working in the autonomous-loop framework until ALL of the following are simultaneously TRUE: - docs/Issues.md Status-column has zero In progress / Ready for testing / In testing / Reopened entries (§11.4.15 closed-set values). - docs/CONTINUATION.md §3 “Active work” section is empty. - No background subagent is

mid-execution (TaskList reports all tasks completed or pending with pending items NOT being prerequisite-blockers for anything actionable). - No external dependency is in-flight (build, flash, push, sync).

If any of those conditions is FALSE, the agent MUST continue working — by claiming the next item from the priority queue, by dispatching a subagent for parallelisable work, or by polling a background task that will be the wake signal. An agent that schedules a wake-up and ends the turn while Issues.md still has actionable items, OR while CONTINUATION.md §3 still has active work, OR while a subagent is still working productively, is in violation.

(B) Zero-idle agent dispatch. When the agent identifies parallelisable work that does not contend on the same file-scope as the main work stream, the agent MUST dispatch a background subagent rather than serialising the work. The main work stream + every background subagent operate concurrently. “Wait for the results of something” is the ONLY acceptable reason for an agent (main or subagent) to be idle. The §11.4.20 + §11.4.58 + §11.4.70 subagent-driven covenants compose with this anchor.

(C) Comprehensive test coverage with real (physical) proofs. Every workable item closed in the loop MUST land four-layer test coverage per §11.4.4(b) — pre-build gate + on-device test + paired §1.1 meta-test mutation + HelixQA Challenge bank entry — and every test PASS MUST cite a captured-evidence artefact per §11.4.5 + §11.4.69. The phrase “physical proofs” is interpreted as: captured audio (tinycap WAV with RMS + fprobe channels), captured video (screen recording + fprobe frames + recording-analyzer event match), captured network state (dumpsys + sink-side probe per §11.4.13), captured UI state (uiautomator dump per §11.4.48), captured sysfs state (/sys/.../proc/... snapshots). Metadata-only PASS, configuration-only PASS, absence-of-error PASS, grep-without-runtime PASS are all critical defects regardless of the green summary line.

(D) Anti-bluff guarantee end-to-end. The §11.4 covenant family (§11.4.1 FAIL-bluffs / §11.4.2 recorded-evidence / §11.4.6 no-guessing / §11.4.7 demotion-evidence / §11.4.27 no-fakes-beyond-unit / §11.4.50 deterministic-consistency / §11.4.52 autonomous-validation / §11.4.68 audio sink-side / §11.4.69 universal sink-side / §11.4.83 docs/qa/ end-user evidence) is the operative truth-discipline of the endless-loop covenant. Every closure narrative MUST cite captured-evidence; every demotion MUST cite same-conditions retest; every status transition MUST cite the deterministic-consistency baseline. **Tests AND Challenges (HelixQA) are bound equally** — a Challenge that scores PASS on a non-functional feature is the same class of defect as a unit test that does. The User mandate explicitly invokes the historical forensic anchor: “we had been in position that all tests do execute with success and all Challenges as well, but in reality the most of the features does not work and can’t be used! This MUST NOT be the case and execution of tests and Challenges MUST guarantee the quality, the completion and full usability by end users of the product!”

(E) Termination criteria. The endless-loop terminates ONLY when: 1. ALL conditions in (A) are simultaneously TRUE (zero open items + zero active CONTINUATION work + zero subagents in flight + zero external dependencies in-flight); OR 2. The operator explicitly issues a STOP / END LOOP / pause autonomous work instruction; OR 3. The host-session-safety covenant (§12 + §12.6 + §12.7 + §12.8 + §12.9 + §12.10) demands suspension to protect the

operator's session; OR 4. The agent has scheduled a wake-up to resume against a known-future-actionable signal (e.g. CI completion, build completion, hardware-attended phase).

No escape hatch. No `--idle-OK`, `--skip-endless-loop`, `--bluff-permitted-for-this-task`, `--metadata-only-test-suffices`, `--no-physical-proof-required` flag exists. The covenant is uniform across every iteration.

Composes with §11.4 (anti-bluff covenant — historical forensic anchor of the project), §11.4.1 (FAIL-bluffs), §11.4.2 (recorded-evidence), §11.4.4 (four-layer coverage), §11.4.5 (audio + video 5-layer quality), §11.4.6 (no-guessing), §11.4.7 (demotion-evidence), §11.4.20 + §11.4.58 + §11.4.70 (subagent-driven defaults), §11.4.27 (no-fakes-beyond-unit + 100% test-type coverage), §11.4.42 (iteration-discipline), §11.4.43 (TDD-fix), §11.4.50 (deterministic-consistency), §11.4.52 (autonomous-validation), §11.4.68 (audio sink-side), §11.4.69 (universal sink-side), §11.4.83 (docs/qa/ end-user evidence), §11.4.85 (stress + chaos), §11.4.86 (roster/corpus auto-sync), §12 family (host-session safety boundaries), §12.10 (CONTINUATION-document maintenance — the source-of-truth state the loop checks against).

Pre-build gate CM-COVENANT-114-87-PROPAGATION enforces this anchor literal in every CLAUDE.md / AGENTS.md / QWEN.md across the canonical fleet (parent + consumer submodules). Paired meta-test mutation strips the literal 11.4.87 → gate FAILs.

Canonical authority: this Constitution.md §11.4.87 in the HelixConstitution submodule ([git@github.com:HelixDevelopment/HelixConstitution.git](https://github.com/HelixDevelopment/HelixConstitution.git)) — inherited per §11.4.35 by every consuming project's repo-root CLAUDE.md / AGENTS.md / QWEN.md.

Non-compliance is a release blocker regardless of context. Skipping the loop, idling without a waited-on signal, or accepting bluff-evidence PASS is severity-equivalent to a §11.4 PASS-bluff at the project-execution layer.

§11.4.88 — Background-push mandate: commit-lock release immediately after commit, push runs detached (User mandate, 2026-05-26)

Forensic anchor — verbatim User mandate (2026-05-26):

“Make sure all these pushes are being done ALWAYS in background in parallel with main work stream so we do not lose time waiting. Everything is committed anyway? If we can do this add this as mandatory rule / constraint that we must respect and follow ALWAYS! ... We MUST ensure that main work stream has always something to do, or wait for the results only when that is absolutely required! We have enormous amount of work in front of us, short deadlines and we MUST DELIVER!”

Forensic incident. 2026-05-26 between 14:35Z and ≥19:30Z (~5 hours observed), a single `commit_all.sh` invocation held its `.git/.commit_all.lock` flock the entire time — the COMMIT had succeeded within seconds at 63d307234ce, but the synchronous `do_push` call kept the flock while `git push vasicdigital_gitlab` did a large initial-mirror upload. **Every subsequent `commit_all.sh` invocation in that window was BLOCKED** — 21 submodule pointer bumps + parent §11.4.87

propagation work + §KA Critical audio quality matrix artifacts all sat staged in the working tree waiting for the push to finish. This is the exact anti-pattern §11.4.87 (zero-idle) prohibits at the project-execution layer.

The mandate.

(A) Flock release IMMEDIATELY after commit lands. Once `git commit` returns 0 (commit object recorded in the local repo), the `.git/.commit_all.lock` MUST be released BEFORE `do_push` is invoked. The commit is durable on local disk regardless of push outcome — there is no correctness reason to gate further local work on a remote-push round-trip.

(B) Push runs detached. After flock release, the push step is spawned via `nohup ./push_all.sh ... > <log> 2>&1 &` then `disown` so the orchestrator's exit does not propagate SIGTERM. The orchestrator's exit code reports COMMIT success, NOT push success (push success is logged separately and reconciled by §11.4.71 `fetch-before-edit` on the next interactive turn).

(C) Push concurrency control: per-remote serialization, multi-remote parallelism. `push_all.sh` MUST acquire a per-remote flock (`.git/.push.<remote>.lock`) so two concurrent invocations targeting the same remote serialize (avoid push-race / non-fast-forward), but invocations targeting DIFFERENT remotes run in parallel. This converts the prior 5-hour serial-stall pattern into a multi-stream pipeline where the operator's main work stream never blocks on the slowest mirror.

(D) Failure surface. Backgrounded push failures land in `qa-results/push_failures/<timestamp>_<remote>.log`. The next interactive autonomous-loop tick MUST check that directory (per §11.4.87(A) “no external dependency is in-flight” check) and surface failures — re-push attempted automatically, or operator notified if auth/quota issue. Silent push-failure is a §11.4 PASS-bluff at the distribution layer.

(E) Synchronous-push escape. An explicit `--sync-push` CLI flag on `commit_all.sh` preserves the legacy synchronous behaviour for `--force-with-lease` paths (per §11.4.41 `force-push merge-first` audit) where the push outcome MUST gate the next step. This is the ONLY escape — without it, every push is backgrounded.

Composes with §2.1 (multi-upstream push norm — §11.4.88 makes it non-blocking), §9.2 (data safety — backgrounded push still uses hardlinked-backup pre-op), §11.4.41 (`force-push merge-first` — uses `--sync-push` escape), §11.4.42 (iteration discipline — endless-loop ticks no longer waste budget on remote-network latency), §11.4.71 (`fetch-before-push merge-first` — applies before background push spawns), §11.4.87 (endless-loop / zero-idle — §11.4.88 is the implementation seam that makes §11.4.87(B) genuinely zero-idle for the dominant blocker class).

Pre-build gate CM-COVENANT-114-88-PROPAGATION enforces this anchor literal in every CLAUDE.md / AGENTS.md / QWEN.md across the canonical fleet. Pre-build gate CM-BACKGROUND-PUSH-WIRED verifies `commit_all.sh` releases flock before `do_push` AND spawns push via `nohup ... &` AND `push_all.sh` acquires per-remote flock. Paired §1.1 meta-test mutations strip the load-bearing literals → gates FAIL.

Canonical authority: this Constitution.md §11.4.88 in the HelixConstitution submodule ([git@github.com:HelixDevelopment/HelixConstitution.git](https://github.com:HelixDevelopment/HelixConstitution.git)).

Non-compliance is a release blocker. Synchronous push (without `--sync-push` flag) is severity-equivalent to a §11.4 PASS-bluff at the project-execution layer. No escape hatch beyond `--sync-push` for §11.4.41 force-push events.

§11.4.89 — Background test execution mandate (User mandate, 2026-05-27)

Forensic anchor — verbatim user mandate (2026-05-27):

“Any tests we are executing, especially long test cycles, MUST BE performed in background in parallel with main work stream! This MUST NOT block our capabilities to work on queued workable items (and tackle main Issues document(s)). Make sure we strictly follow this starting now! Main work stream can be blocked or sit iddle only if absolutely needed and if it depends hard on results of some background execution.”

Forensic incident. 2026-05-27 the conductor invoked `bash device/rockchip/rk3588/tests/pre_build_verification.sh` synchronously with a timeout 360 foreground wrapper — even though the operator had previously mandated §11.4.87 endless-loop zero-idle. The 6-7 minute test run blocked the entire main work stream from progressing on §JV/§JW/§JX/§JY scaffolding + further Issues.md drains. This is the exact anti-pattern §11.4.87(B) prohibits at the test-execution layer (analogous to §11.4.88 at the push-execution layer).

The mandate.

(A) Long-running tests run detached. Any test cycle expected to exceed ~30 seconds — `pre_build_verification.sh`, `meta_test_false_positive_proof.sh`, `test_all_fixes.sh`, `recent_work_validate.sh`, HelixQA Challenge banks, on-device 4-phase cycles, full-suite retests per §11.4.40, `audio_loop_supervisor.sh`, `dual_display_record.sh` harnesses, `verify-all-constitution-rules.sh` — MUST be spawned via `nohup ... > <log> 2>&1 & + disown` and `<log>` placed under a known directory (e.g. `qa-results/<test_id>_<timestamp>.log`) for later inspection.

(B) Main stream proceeds on queued items. The conductor MUST immediately return to the §11.4.42 priority queue (open Issues.md items, queued PWUs, doc-sync work, audit work, anything not depending hard on the in-flight test’s exit code). “Wait for results” is the ONLY acceptable idle reason per §11.4.87(B) — and even then, the conductor checks back on the test result via short polling (`tail / mtime / exit-status` file) rather than blocking on wait.

(C) Hard-dependency gating. A subsequent step that genuinely requires the test’s exit code (e.g. `commit_all.sh` refuses if meta-test is in flight per §11.4.84; release-tag gate per §11.4.40 step 4) MUST express that dependency explicitly — poll the exit-status file or check `pgrep -af <test-name>` before proceeding. If the test is still running and the next step needs its result, surface that to the operator per §11.4.66 with interactive options.

(D) Failure surface + cleanup. Backgrounded test failures land in their <log> files. The next autonomous-loop tick MUST check that directory (analogue of §11.4.88(D)). Test-script-level cleanup (per §11.4.14 + §11.4.84 + this anchor's restore discipline) is non-negotiable — `trap '<cleanup>' EXIT INT TERM` in every long-running test that mutates global state (mid-mutation backup files, sysfs nodes, ALSA states, etc.).

(E) Pre-empt the blocking pattern. The conductor MUST NOT invoke `bash <long-test>.sh` synchronously when a backgrounded variant suffices. Foreground execution is permitted ONLY when: (a) the test completes in <30 s, OR (b) the operator explicitly requests foreground execution for a specific step. Synchronous execution of a multi-minute test absent operator authorisation is a §11.4.89 violation regardless of how much “context-cache benefit” the foreground path appears to offer.

(F) Concurrency control. Multiple backgrounded tests of the SAME script MUST serialise via per-script flock (e.g. `.git/.pre_build.lock`, `.git/.meta_test.lock`). Concurrent invocations of DIFFERENT tests run in parallel.

Composes with §11.4.42 (iteration discipline — backgrounding tests recovers the operator's wait-time per §11.4.82(F) parallel multi-device testing), §11.4.66 (interactive clarification — surface hard-dependencies as options instead of blocking), §11.4.82 (iteration-speedup discipline — §11.4.89 is the dominant blocker class at the test layer, analogous to §11.4.88 at the push layer), §11.4.84 (working-tree quiescence — backgrounded mutation tests still need quiescence checks), §11.4.85 (stress + chaos — long stress runs are quintessential §11.4.89 candidates), §11.4.87 (endless-loop zero-idle — §11.4.89 is the implementation seam at the test-execution layer that makes §11.4.87(B) genuinely zero-idle for the test blocker class), §11.4.88 (background-push — §11.4.89 is the symmetric anchor for tests).

Pre-build gate CM-COVENANT-114-89-PROPAGATION enforces this anchor literal across the canonical fleet. Pre-build gate CM-BACKGROUND-TEST-EXECUTION-WIRED verifies that long-running test invocations in helper scripts use `nohup ... & + disown` pattern (or equivalent). Paired §1.1 meta-test mutations strip the load-bearing literals → gates FAIL.

Canonical authority: this Constitution.md §11.4.89 in the HelixConstitution submodule (`git@github.com:HelixDevelopment/HelixConstitution.git`).

Non-compliance is a release blocker. Synchronous long-test execution (without operator authorisation) is severity-equivalent to a §11.4 PASS-bluff at the project-execution layer. No escape hatch beyond explicit per-invocation operator authorisation.

§11.4.90 — Obsolete status + per-item obsolescence audit mandate (User mandate, 2026-05-27)

Forensic anchor — verbatim user mandate (2026-05-27):

“Regarding the Critical Bug No 6 - Albums cover etc. - Now it seems obsolete after latest request for new behavior when audio and video content is being played. If we do not have Obsolete as an option / status for the ticket in its resolving, we should introduce this now and mark obsolete tickets with some light gray background or something that will indicate that item is no longer valid. Maybe text - the description to be strikethrough styled as well! Once all this is done please review all existing open or resolved workable items (fixed as well) if they are obsolete - not valid any more. ... There MUST NOT be any mistake! No bluff is allowed of any kind!”

The §11.4.15 Status closed-set is extended with a 4th terminal value **Obsolete** (→ **Fixed.md**) matching the existing 3-element §11.4.33 closure-vocabulary table (Bug → Fixed, Feature → Implemented, Task → Completed). The **Obsolete** value applies regardless of Type when the workable item is **no longer valid** because: (a) the underlying behaviour the item described has been superseded by a later operator mandate / design pivot (canonical example: Bug #6 Albums cover, superseded by §JU + the 2026-05-27 secondary-display video-playback redesign), (b) the affected feature was REMOVED from the product, (c) the item was duplicated and the canonical entry resolved it, (d) the item was filed against a hardware/topology variant we no longer support.

Mandatory obsolescence audit metadata — every **Obsolete** (→ **Fixed.md**) heading **MUST** carry, within 8 non-blank lines of the heading, an ****Obsolete-Details:**** line capturing four sub-facts (mirroring §11.4.21 + §11.4.34 audit-line patterns):

- **Since:** ISO date (YYYY-MM-DD).
- **Reason:** one-line cause from the closed vocabulary { superseded-by-design-change | superseded-by-later-mandate | feature-removed | duplicate-of | unsupported-topology }.
- **Superseding-item:** §-letter reference or User-mandate verbatim quote anchor citing the work that obsoleted it.
- **Triple-check evidence:** path to captured evidence (git log, code grep, runtime behaviour) confirming the item is genuinely no longer valid. Per operator mandate “There MUST NOT be any mistake” — bare assertion is forbidden; positive captured evidence per §11.4.6 is mandatory.

Visual treatment — the §11.4.23 colorizer **MUST** style **Obsolete** cells with `cell-status-obsolete` class: light-gray background #E0E0E0 + strikethrough text on the row’s description cell. Operator mandate verbatim: “mark obsolete tickets with some light gray background or something that will indicate that item is no longer valid. Maybe text - the description to be strikethrough styled as well”.

Audit cadence — at every release-gate sweep (per §11.4.40), the conductor **MUST** re-evaluate every non-terminal **Issues.md** item **AND** every **Fixed.md** item for obsolescence per the closed-set Reason vocabulary above. New obsolescence findings produce **Obsolete** (→ **Fixed.md**) migrations atomic per §11.4.19. Triple-check is non-negotiable: each obsolescence claim cites (i) the superseding mandate, (ii) code-state grep confirming the obsoleted behaviour is gone, (iii) runtime/test evidence confirming no test still exercises the obsoleted path.

Pre-build gates **CM-COVENANT-114-90-PROPAGATION** (anchor literal across canonical files) + **CM-ITEM-OBSOLETE-DETAILS** (every **Obsolete** heading carries ****Obsolete-Details:**** line within 8 lines) + **CM-OBSOLETE-COLORIZER-WIRED**

(§11.4.23 colorizer emits `cell-status-obsolete` class on the right cells + the CSS file defines the class). Paired §1.1 mutations strip the anchor literal / delete an Obsolete-Details line / strip the colorizer class — every mutation FAILs its gate.

Composes with §11.4.15 (item-status closed-set extension), §11.4.16 (Type tracking unchanged — Obsolete is orthogonal to Type), §11.4.19 (Fixed-document column-alignment — Obsolete migrates to Fixed.md just like Fixed/Implemented/Completed), §11.4.21 (Operator-blocked details-line pattern — Obsolete-Details mirrors), §11.4.23 (visual cue + grouping — adds `cell-status-obsolete`), §11.4.33 (closure vocabulary — Obsolete is the 4th terminal value), §11.4.34 (Reopened-source pattern — Obsolete uses analogous audit line), §11.4.40 (release-gate sweep — audit cadence trigger), §11.4.42 (iteration discipline — obsolescence audit is a routine cycle step), §11.4.66 (interactive clarification — surface ambiguous obsolescence to operator), §11.4.71 (pre-push fetch — re-check obsolescence post-fetch).

Canonical authority: this Constitution.md §11.4.90 in the HelixConstitution submodule.

Non-compliance is a release blocker. Marking a still-valid item Obsolete is a §11.4 PASS-bluff at the planning layer. Marking a genuinely-obsolete item with a non-terminal Status is documentation drift. No escape hatch — every obsolescence finding follows the triple-check evidence rule above.

§11.4.91 — Summary-doc clarity mandate (User mandate, 2026-05-27)

Forensic anchor — verbatim user mandate (2026-05-27):

“We see in Summary docs - Issues_Summary some not clear one line descriptions - like for example: ‘Composes with’. For each workable item in any variant of documentation - long or short descriptions - we **MUST HAVE** clearly understandable meaning of particular entry of the document! Re-evaluate all of them, fix the descriptions (especially one liners) and make sure that every team member can clearly understand what that particular or any particular workable item from the list is exactly about! There cannot be misunderstanding or unclarity of any kind and no bluff allowed!”

Every short-form summary entry (Issues_Summary.md, Fixed_Summary.md, README.md doc-link section, Status_Summary.md page 1 + 2, every other one-liner derived from a long-form tracker entry) **MUST** contain a description that is **self-contained + meaningfully informative** about **WHAT** the item is, not a fragment of the long-form body. Specifically forbidden one-liner anti-patterns:

- **Section labels** as descriptions: Composes with, Closure criteria, Fix direction, Forensic anchor, Reopened-Details, Operator-Block-Details, Obsolete-Details, Summary, Status, Type, Severity, etc. — these are section *labels* inside the long-form entry, **NEVER** appropriate as the one-liner description.
- **Bare numeric / categorical fragments:** Critical, High, Medium, Low, Bug, Feature, Task, Fixed, standalone — these are metadata column values, not descriptions.

- **Generic restatements of Status:** In progress, Queued, Reopened, Operator-blocked, Fixed, standalone — again metadata column values.
- **Section-marker echoes:** anything matching `^[A-Z][a-z]+ (with|by|on|in|to|for):?$` patterns where the right-hand side is empty.
- **§-letter alone:** §KB, §EU, etc. — the section letter is the identifier, not the description.

Required one-liner shape — every summary entry’s description column **MUST** contain a complete clause (≥ 6 words OR ≥ 40 characters, whichever is longer) that names the SUBJECT (component / app / behaviour / defect class) + the PROBLEM or GOAL (what’s broken / what’s being added / what’s being audited). Example transformations:

- Bad: “Composes with”
- Good: “MPV fork comprehensive secondary-display playback + interaction-pattern + chaos test suite (§KB-2)”
- Bad: “Critical”
- Good: “HDMI audio total loss with ‘command error 1’ on D3 post-reflash (Tier 1/2/3 source-side fixes landed, REQUIRES_REBUILD validation pending)”
- Bad: “§JU”
- Good: “Presenter 2nd-display behaviour simplification + wake-on-play + notification dim toggle (umbrella, REQUIRES_REBUILD pending validation)”

Generator-level enforcement — `generate_issues_summary.sh` + `generate_fixed_summary.sh` + `update_readme_doc_links.sh` + `generate_status_summary.sh` **MUST** extract the one-liner from the **H1/H2 heading line** of the source long-form entry (which already encodes the subject + context per §11.4.15 + §11.4.16 + §11.4.54 ATM-NNN convention), **NEVER** from arbitrary downstream text. The generators **MUST** refuse to emit a row whose description matches any of the forbidden anti-patterns above — emitting (MISSING DESCRIPTION – fix source heading) placeholder instead, with the offending row visually highlighted.

Pre-build gate CM-SUMMARY-CLARITY-DESCRIPTIONS scans every `*_Summary.md` for the forbidden anti-patterns above; any match **FAILs** the gate. Paired §1.1 mutation injects “Composes with” into a Summary cell → gate **FAILs**.

Audit cadence — at every release-gate sweep (§11.4.40 + §11.4.42 step 4), re-evaluate every summary row for clarity. Operator-facing exports (HTML + PDF) **MUST** never ship with an anti-pattern row.

Composes with §11.4.12 (Issues_Summary sync — clarity is a property of the generated row), §11.4.19 (Fixed-document column alignment — clarity applies symmetrically), §11.4.23 (colorizer — anti-pattern rows highlighted), §11.4.44 (revision header — generator updates revision on clarity regeneration), §11.4.53 (Fixed_Summary parity), §11.4.56 (Status_Summary two-audience format — page 1 audience **MUST** get the clearest descriptions), §11.4.57 (README doc-link rows — same anti-pattern rule), §11.4.59 (README always-sync), §11.4.60 (composite docs always-sync), §11.4.65 (universal MD export), §11.4.74 (mechanical enforcement — pre-commit hook for summary clarity).

Canonical authority: this Constitution.md §11.4.91 in the HelixConstitution submodule.

Non-compliance is a release blocker. A summary row whose description is an anti-pattern fragment is severity-equivalent to a §11.4 PASS-bluff at the documentation layer — it implies the item is documented when it is not. No escape hatch.

§11.4.92 — Multi-pass change-evaluation discipline (User mandate, 2026-05-27)

Forensic anchor — verbatim user mandate (2026-05-27):

“Every change to the project or codebase we do MUST BE evaluated in several passes and in in-depth analysis for potential new issues or problems it can introduce! We MUST BE sure that: change nothing breaks, we do validate and verify this in codebase and through comprehensive deep research, that no new issues or problems of any kind is bringing and that main task is achieved with particular change with no bluff of any kind! After we do change or set of changes this mandatory steps MUST BE taken! Then, and only then, when we have validate and verified codebase and performed full analysis and taken into the account all existing knowledge(s) and information sources and confirmed all points we have mentioned now, then this can be accepted as change that is going to be committed, pushed, tested, released and so on! Write as many as required additional documentation and notes about our codebase, the system, or the changes so we can achieve this goals!”

Every non-trivial change to the project codebase MUST pass a multi-pass evaluation BEFORE the change is accepted as commit-ready. The discipline expands §11.4.4 (test-interrupt-on-discovery) + §11.4.8 (deep-web-research) + §11.4.43 (TDD-fix) + §11.4.50 (deterministic consistency) + §11.4.82 (Phase 1 forensic before any speculative patch) into an explicit pre-commit evaluation pipeline.

The 5-pass evaluation contract (mandatory, ALL passes must hold):

Pass 1 — Main-task verification. The change achieves the stated main task. Captured-evidence per §11.4.5 / §11.4.69 demonstrates the user-visible behaviour now works AS INTENDED. NO inferred / NO assumed / NO “should work” — only captured evidence.

Pass 2 — Regression-blast-radius analysis. Every file the change touches AND every file that imports / sources / references the changed file is enumerated. For each, the agent demonstrates (via grep / dependency-graph / actual test run) that the change does not break the existing contract. Examples: function signature changes → audit every caller; HAL config change → audit every test consuming the config; gate addition → audit every paired meta-test mutation. Captured evidence per §11.4.5.

Pass 3 — Cross-feature interaction analysis. Beyond direct-dependency blast radius, the change is evaluated against parallel features that share state, timing, hardware, or shell environment. Examples: audio HAL change → check video

routing (Fix #88 SurfaceView Z-order timing); WiFi roaming change → check BT A2DP coex (Fix #70); test-script addition → check §11.4.84 working-tree quiescence + §11.4.89 background-test discipline. Captured evidence per §11.4.5.

Pass 4 — Deep-research validation. Per §11.4.8, the chosen approach is verified against external state-of-the-art (official docs, vendor guides, open-source codebases, kernel mailing list, AOSP gerrit, GitHub issues). Either the change cites an external precedent OR the literal “NO external solution found — original work” is recorded. CodeGraph queries per §11.4.78 + §11.4.79 surface intra-codebase precedents.

Pass 5 — Anti-bluff confirmation. Per §11.4 / §11.4.1 / §11.4.6 / §11.4.27 / §11.4.50 / §11.4.52 / §11.4.69 / §11.4.83 the change is verified to NOT introduce any new bluff surface — no metadata-only PASS, no config-only PASS, no script-bug FAIL-bluff, no “tests pass but feature doesn’t work” anti-pattern. Captured evidence directories per §11.4.5 reference the exact passing paths.

Documentation requirement. Each Pass produces written documentation (commit-message footers OR docs/ entries OR qa-results/ evidence) demonstrating the pass was completed. A change without all 5 passes documented is severity-equivalent to a §11.4 PASS-bluff at the development-process layer.

Acceptance criteria. Only AFTER all 5 passes complete with documented evidence may the change be: (a) committed via `commit_all.sh`, (b) pushed via §11.4.88 background-push, (c) tested via §11.4.89 background test, (d) released / tagged per §11.4.40.

No escape hatch. No `--skip-multi-pass`, `--single-pass`, `--bluff-permitted`, `--fast-commit` flag exists. The 5 passes apply to every non-trivial change per the operator’s verbatim mandate. Trivial changes (typo fixes, doc revision-header bumps, MD-export regeneration) are exempt ONLY when (i) the change touches zero source code AND (ii) the commit message explicitly cites the exemption.

Composes with §11.4 / §11.4.1 (anti-bluff baseline), §11.4.4 (test-interrupt-on-discovery — Pass 2 + Pass 3 trigger this), §11.4.5 (captured-evidence quality — every Pass’s evidence MUST be quality-checked), §11.4.6 (no-guessing — Pass 1-3 require captured evidence not hypothesis), §11.4.8 (deep-research — Pass 4 expansion), §11.4.20 / §11.4.70 (subagent-driven — Pass 2-3 dispatchable to subagents per parallel analysis), §11.4.27 (no-fakes-beyond-unit — Pass 5 enforces), §11.4.42 (iteration discipline — multi-pass IS the iteration), §11.4.43 (TDD-fix — RED→Pass-1, GREEN→Pass-2-5), §11.4.50 (deterministic consistency — Pass 1 captured evidence MUST be N=3 reproducible), §11.4.52 (autonomous validation — Pass 5 anti-bluff confirmation gates on autonomous evidence), §11.4.69 (universal sink-side evidence — Pass 1+5 use closed-set taxonomy), §11.4.78 / §11.4.79 (CodeGraph — Pass 3 dependency-graph queries), §11.4.82 (Phase 1 forensic — §11.4.92 generalises Phase 1 to Pass-1+Pass-2+Pass-3 across every change class), §11.4.83 (docs/qa transcript — Pass evidence dir is the QA transcript), §11.4.85 (stress + chaos — Pass 2+3 evidence MUST include stress + chaos signals when applicable), §11.4.87 (endless-loop — multi-pass runs concurrently across PWUs per §11.4.58), §11.4.89 (background tests — Pass 1-3 long tests run backgrounded).

Pre-build gate CM-COVENANT-114-92-PROPAGATION enforces this anchor literal across the canonical fleet. Pre-build gate CM-MULTI-PASS-EVALUATION-EVIDENCE audits recent commits for the 5-pass evidence trail (commit message footer OR per-commit qa-results/<commit-hash>/passes/ directory OR Issues.md / Fixed.md narrative). Paired §1.1 meta-test mutations strip the load-bearing literals → gates FAIL.

Canonical authority: this Constitution.md §11.4.92 in the HelixConstitution submodule.

Non-compliance is a release blocker. A commit landing without the 5-pass evaluation documented is severity-equivalent to a §11.4 PASS-bluff at the development-process layer — it implies the change was vetted when it was not.

§11.4.93 — SQLite-backed single-source-of-truth for workable items (User mandate, 2026-05-27)

Forensic anchor — verbatim user mandate (2026-05-27):

“There MUST be single source of truth for all of our workable items - SQLite database containing all workable items, all data about it and from there proper scripts (we recommend Go programs) which will update / regenerate all documentation - Issues, Fixed, summary docs, and everything related! We MUST reduce a chance for sync to be broken between the Db and all documents we have! We MUST BE able to generate always all docs from DB or to re-generate Db from all docs we have in opposite direction if we must! All this MUST BE covered with all supported test types, validated and verified by tests which will produce real proofs and be 100% anti-bluff compatible!”

The current text-based Issues.md / Fixed.md / Issues_Summary.md / Fixed_Summary.md / CONTINUATION.md / Status.md tracker constellation is converted to a **SQLite-database-backed single source of truth** with bidirectional regeneration between DB and Markdown. The DB is the authoritative source; all Markdown / HTML / PDF / Status / Summary surfaces are generator output. Sync drift is mechanically impossible because every regeneration starts from the DB.

Canonical artefacts (all consuming projects MUST adopt):

1. **Database file** at canonical path docs/workable_items.db (SQLite 3, **TRACKED in git, NEVER gitignored** per User mandate 2026-05-27 amendment). The DB IS the single source of truth — gitignoring it would defeat the SSoT mandate at the version-control layer. Every commit that mutates workable items MUST stage + push the DB file alongside the Markdown regen output. The DB is an exception to §11.4.30 build-artefact gitignore — it is NOT a build artefact, it IS authoritative source data. SQLite’s WAL mode produces a sidecar *.db-wal + *.db-shm which ARE gitignored (transient WAL files; their state is checkpointed into the main .db file by workable-items sync or any explicit PRAGMA wal_checkpoint(TRUNCATE)).

2. Schema (mandatory minimum tables):

- `items` — primary key `atm_id` per §11.4.54; columns type (§11.4.16 closed-set), status (§11.4.15 + §11.4.90 closed-set including `Obsolete`), severity, title, description (≥ 6 words / ≥ 40 chars per §11.4.91), `created_at`, `last_modified`, `current_location` (`Issues|Fixed`), `forensic_anchor`, `closure_criteria`, `composes_with` (JSON array of refs).
- `item_history` — append-only audit log: `atm_id`, `event_type` (`Opened|Updated|Reopened|Fixed|Implemented|Completed|Obsolete`), `by` (`AI|User` per §11.4.34), `on_date`, `reason` (closed vocabulary per §11.4.34 + §11.4.90), `evidence_path`.
- `obsolete_details` — `atm_id`, `since`, `reason` (§11.4.90 closed-set), `superseding_item`, `triple_check_evidence` (mandatory per §11.4.90).
- `operator_block_details` — `atm_id`, `what`, `why_exhausted_alternatives`, `unblock_condition`, `who` per §11.4.21.
- `firebase_metadata` — per §11.4.47.
- `meta` — schema version, `last_sync_direction`, `last_sync_timestamp`, `integrity_hash`.

3. Go binary at canonical path `cmd/workable-items/` (separate Go module, NOT part of AOSP build). Subcommands:

- `workable-items sync md-to-db` — parse `Issues.md` + `Fixed.md` (single source of authority during transition window), upsert into DB. Idempotent.
- `workable-items sync db-to-md` — regenerate `Issues.md` + `Fixed.md` + `Issues_Summary.md` + `Fixed_Summary.md` + `CONTINUATION.md` §3 + every `Status.md` from DB. Idempotent.
- `workable-items diff` — print DB vs MD divergence. Used by pre-build gate.
- `workable-items validate` — schema sanity + §11.4.15/§11.4.16/§11.4.33/§11.4.34/§11.4.54/§11.4.90/§11.4.91 invariants per row. FAIL on any violation.
- `workable-items add <type> <severity> --title <title> --description <description>` — interactive entry with mandatory fields enforced.
- `workable-items close <atm-id> --status <fixed|implemented|completed|obsolete> --evidence <path>` — terminal transition with mandatory captured-evidence per §11.4.5/§11.4.90.

4. Operational integration:

- `commit_all.sh` pre-commit hook runs `workable-items diff` and refuses if MD/DB diverge.
- `sync_issues_docs.sh` invokes `workable-items sync db-to-md` instead of legacy bash generators (`generate_issues_summary.sh` becomes a shim wrapping the Go binary).
- `pre_build_verification.sh` runs `workable-items validate` as a gate; failure increments `ERRORS`.

5. Anti-bluff test coverage (mandatory per §11.4 + §11.4.85 + §11.4.27):

- Unit tests for every CRUD operation + every schema invariant.

- Integration test: full round-trip MD→DB→MD with byte-identical re-emission (modulo cosmetic-whitespace normalisation per closed-set tolerance).
- Stress + chaos per §11.4.85: 1000-row insert / concurrent-write (10 writers) / mid-write SIGKILL / corrupt-DB recovery / disk-full handling.
- Paired §1.1 meta-test mutation: strip a CRUD function → unit test FAILs.
- HelixQA Challenge entry CME-WORKABLE-ITEMS-001 exercising end-to-end DB→MD→DB→MD round-trip.

Bidirectional regeneration guarantee. The DB MUST be reconstructible from the MD docs (for disaster recovery + audit replay) AND the MD docs MUST be reconstructible from the DB (the primary direction during normal operations). Either direction's regeneration produces semantically equivalent output (closed-set tolerance for cosmetic whitespace + section-ordering noise).

Cross-project propagation. Every consuming project that adopts this constitution submodule inherits the SQLite-SSoT approach. The Go binary lives in the constitution submodule (constitution/scripts/workable-items/) so consumers reference it from there per §11.4.74 catalogue-first discipline — never reimplement.

Migration path (per §11.4.42 iteration discipline + §11.4.58 PWU pipeline): - **Phase 1:** file §LA Issues entry tracking the migration, scope-locked. - **Phase 2:** Go binary scaffold + schema DDL committed. - **Phase 3:** sync md-to-db lands + initial migration captures current state. - **Phase 4:** sync db-to-md lands + byte-identical round-trip CI gate. - **Phase 5:** existing generators (generate_issues_summary.sh etc.) become Go-binary shims. - **Phase 6:** legacy text-direct edits prohibited (pre-commit hook).

Pre-build gates CM-COVENANT-114-93-PROPAGATION (anchor literal in canonical fleet) + CM-WORKABLE-ITEMS-DB-PRESENT (DB regen mechanism per §11.4.77 + DB schema-version current) + CM-WORKABLE-ITEMS-MD-DB-IN-SYNC (diff returns empty). Paired §1.1 meta-test mutations strip the load-bearing literals → gates FAIL.

Composes with §11.4 (anti-bluff covenant — DB is the captured-evidence registry), §11.4.12 (Issues_Summary sync — DB-driven), §11.4.15 (Status closed-set — DB schema enforces), §11.4.16 (Type closed-set — DB schema enforces), §11.4.17 (universal-vs-project — DB schema is universal), §11.4.19 (Fixed-document column-alignment — DB query emits both Summaries from same query), §11.4.21 (Operator-blocked details — operator_block_details table), §11.4.27 (no-fakes-beyond-unit — DB integration tests exercise real SQLite), §11.4.30 (.gitignore + regeneration mechanism — DB file gitignored with workable-items sync as regen), §11.4.33 (closure vocabulary — DB enforces), §11.4.34 (Reopened source attribution — DB item_history table), §11.4.42 (iteration discipline — 6-phase migration), §11.4.43 (TDD — RED before each phase), §11.4.44 (revision header — Markdown output preserves), §11.4.45 (Status.md per integration — generator queries DB by domain), §11.4.50 (deterministic consistency — round-trip MD→DB→MD byte-identical), §11.4.51 (LIVE_ADB_FIRST — N/A, host-only), §11.4.52 (autonomous validation — DB integrity is autonomously verifiable), §11.4.53 (Fixed_Summary parity — DB ensures), §11.4.54 (ATM-NNN — DB primary key), §11.4.55 (Reopens-history doc — DB item_history queryable), §11.4.56 (Status_Summary parity — DB queries produce both pages), §11.4.57 (README doc-link section — DB drives), §11.4.58 (parallel PWU — Phase 1-6 land as separate PWUs), §11.4.60 (composite always-sync — DB→MD→HTML+PDF), §11.4.65 (universal MD export — final stage), §11.4.74

(catalogue-first — Go binary in constitution), §11.4.83 (docs/qa transcript — DB queries serve as audit replay), §11.4.85 (stress + chaos — DB integration test scope), §11.4.86 (roster/corpus auto-sync — DB schema accommodates roster tables), §11.4.87 (endless-loop — DB queries are zero-idle-friendly), §11.4.89 (background tests — DB validation runs detached), §11.4.90 (Obsolete status — `obsolete_details` table), §11.4.91 (Summary clarity — DB enforces description-floor at insert time), §11.4.92 (multi-pass evaluation — every DB schema migration follows 5-pass).

Canonical authority: this Constitution.md §11.4.93 in the HelixConstitution submodule.

Non-compliance is a release blocker. A consuming project that maintains text-based trackers without the DB-SSoT migration plan filed + Phase progression actively in flight is severity-equivalent to a §11.4 PASS-bluff at the data-architecture layer — the sync-drift surface §11.4.93 closes IS a §11.4 violation vector.

§11.4.94 — Zero-idle priority-first parallel-by-default operating mode (User mandate, 2026-05-27)

Forensic anchor — verbatim user mandate (2026-05-27):

“We MUST NEVER sit iddle / wait or sleep if there is possibility for us to work on something like in this case! Continue all work further in parallel if possible fully autonomously! We MUST not waste time since work scope we have is huge! ... Always check if there is a possibility to work on something while we are not working actively on something! Pick always by priority - most critical workable items and other tasks MUST BE done first! ... Everything we can should be done fully in background, should not affect stability of the System or create problems, should be done using subagents-driven approach! ... Stay still / iddle if nothing is left to be done at all or waiting for something that is blocking us / you!!!”

§11.4.94 is the **operating-mode reaffirmation** that binds §11.4.87 (endless-loop), §11.4.20 / §11.4.70 (subagent-driven), §11.4.42 (iteration-priority), §11.4.58 (parallel PWU), §11.4.72 (audio top-priority), §11.4.82 (iteration-speedup), §11.4.88 (background-push), §11.4.89 (background-test) into a SINGLE always-on enforcement contract:

The mandate (binding on the conductor + every subagent + every helper script):

(A) Idle-only-when-genuinely-blocked. The conductor MAY sleep / wait / be idle ONLY when (i) every priority-queued workable item is genuinely blocked on external dependency (operator hardware action, network upstream, hardware-attended forensic, build/test completion that the conductor cannot accelerate), OR (ii) the operator has issued an explicit STOP, OR (iii) host-session-safety demands it per §12. “I don’t see what to do next” is NEVER a valid idle reason — the conductor MUST recheck the priority queue per §11.4.42 step 1 and §11.4.87(A) termination criteria.

(B) Always check parallel-work feasibility at every pause point. Before any wake-up schedule / before any “waiting for X” / before any sleep call, the conductor MUST: (1) survey the priority queue per §11.4.42 SCOPE LOCK + §11.4.72 audio-first + §11.4.87(B) zero-idle, (2) identify ALL items that can progress without contending with the in-flight blocker, (3) dispatch them in parallel per §11.4.20 / §11.4.70 (subagent-driven if non-trivial) + §11.4.58 (PWU disjoint scope) + §11.4.89 (background long tests). Only AFTER step (3) returns “no parallel work available” may the conductor schedule a wake/sleep.

(C) Priority order MANDATORY at every pick. Per §11.4.42 priority axis + §11.4.72 audio-top-priority axis. Pick the highest-Severity / highest-priority item the conductor can autonomously progress. Lower-priority items wait until higher-priority queues drain or block on external dependency.

(D) Subagent-driven by default for non-trivial scope. Per §11.4.20 + §11.4.70. The conductor remains the integration + commit + push seam (per §11.4.58 + §11.4.84 quiescence + §11.4.88); subagents execute the analyses + scaffolds + tests + migrations in parallel. Inline execution permitted only when the task is trivial (single-file edit, sub-300-line diff) OR the operator explicitly requested it.

(E) Background by default for long-running work. Per §11.4.85 stress + §11.4.89 background test + §11.4.88 background push. Any operation expected to exceed ~30 s wall-clock MUST run detached (nohup + disown) with log under qa-results/<op-id>_<ts>.log. The conductor returns to the priority queue immediately and polls back when needed.

(F) Stability-preserving — no compromise. Parallel work MUST NOT compromise system stability or create new problems. Composes with §11.4.92 multi-pass change-evaluation — every parallel branch’s output passes 5-pass before integration. Composes with §11.4.84 working-tree quiescence — concurrent subagents in same checkout coordinate via lockfile or git worktree. Composes with §12.6 / §12.7 / §12.8 / §12.9 host-session safety — parallel work respects 60% memory ceiling and -j2 AOSP cap.

(G) Status updates as work progresses. Per operator mandate “Let us know how goes this mandatory approach”. The conductor surfaces parallel-work catalogue + in-flight branches + recent closures in compact summary form whenever the operator asks “how are we progressing” OR at natural milestone boundaries (after each constitution-anchor landing, after each commit batch, after each subagent return).

Anti-pattern explicitly forbidden:

- Scheduling a wake-up without first surveying the parallel-work queue per (B).
- Returning “nothing to do — sleeping” when non-trivial priority items remain that the conductor could autonomously progress.
- Serialising work that could run in parallel without justification.
- Picking lower-priority items while higher-priority ones remain progressable.
- Doing work foreground that should run via subagent or background.
- Letting parallel work introduce instability or regression.

Pre-build gate CM-COVENANT-114-94-PROPAGATION enforces this anchor literal across the canonical fleet. Pre-build gate CM-PARALLEL-WORK-AUDIT (when implemented) audits the last N wake-up schedule decisions for surfaced parallel-work-queue evidence — wake-ups without the survey are §11.4.94 violations. Paired §1.1 meta-test mutations strip the load-bearing literals → gates FAIL.

Composes with §11.4.20 (subagent-driven) / §11.4.42 (iteration discipline) / §11.4.58 (parallel-development PWU) / §11.4.70 (subagent-driven default) / §11.4.72 (audio top-priority) / §11.4.82 (iteration-speedup) / §11.4.84 (working-tree quiescence) / §11.4.85 (stress + chaos) / §11.4.87 (endless-loop zero-idle) / §11.4.88 (background-push) / §11.4.89 (background-test) / §11.4.92 (multi-pass evaluation) / §12.6 + §12.7 + §12.8 + §12.9 (host-session safety).

Canonical authority: this Constitution.md §11.4.94 in the HelixConstitution submodule.

Non-compliance is a release blocker. Conductor idle without exhausting the parallel-work queue is severity-equivalent to a §11.4 PASS-bluff at the project-execution layer — it implies the conductor has progressed the project when it has not. No escape hatch beyond the (A) genuinely-blocked / explicit-STOP / host-safety triad.

§11.4.95 — Workable-items SQLite DB is TRACKED in git, NEVER gitignored (User mandate, 2026-05-27)

Forensic anchor — verbatim user mandate (2026-05-27):

“We shall not Git ignore our workable items SQLite DB since it is our single source of truth for current working items we have and all data related to it! We MUST fix this if DB is Git ignored! ... workable items SQLite DB regularly committed and pushed to all upstreams!”

The §11.4.93 workable-items SQLite database at docs/workable_items.db is the **AUTHORITATIVE source of truth** for every workable item in every consuming project. It MUST be:

- **TRACKED in git** at canonical path docs/workable_items.db. NEVER gitignored regardless of file-size or “build-artefact-class” heuristics — the DB is NOT a build artefact, it IS authoritative source data.
- **Committed alongside every workable-item state change.** Every workable-items sync md-to-db invocation that mutates DB state MUST stage + commit + push the DB file in the same commit as the corresponding MD changes per §11.4.19 atomic-move discipline.
- **Pushed to every upstream per §2.1** — every mirror of every consuming project carries the latest DB so a fresh clone from any mirror reconstructs the full workable-items inventory without consulting any other source.
- **WAL-checkpointed before commit** — SQLite’s WAL mode produces transient *.db-wal + *.db-shm sidecars. Those ARE gitignored (per §11.4.30 transient/cache classification). The workable-items Go binary MUST execute PRAGMA wal_checkpoint(TRUNCATE) before any commit-stage step so the main .db file carries the authoritative state and the sidecars are safely discardable.

- **NEVER force-rewritten** without §9.2 hardlinked-backup + operator authorization. The DB's commit history IS the workable-items audit trail; rewriting it is data-loss equivalent to force-pushing a tracked Issues.md history.

§11.4.30 carve-out. This anchor is an explicit named exception to §11.4.30's build-artefact / data-file gitignore rule. The DB file at docs/workable_items.db is **NOT** a build artefact and is **NOT** transient data — it is the project's SSoT registry whose value increases monotonically as work lands. The §11.4.77 regeneration mechanism applies to other gitignored bulk data (.git-backup-*, RKTtools/linux/, etc.) — NOT to this DB.

§11.4.93 amendment. §11.4.93's earlier text describing the DB as "gitignored per §11.4.30 with §11.4.77 regeneration mechanism" is hereby **AMENDED**: the DB is TRACKED, not gitignored. The Markdown trackers (Issues.md / Fixed.md / Summaries / Status.md fleet / CONTINUATION.md) remain tracked AND derived; the DB is tracked AND authoritative. Both directions of regeneration (workable-items sync db-to-md + workable-items sync md-to-db) continue to land per §11.4.93's bidirectional round-trip guarantee.

Pre-build gates CM-COVENANT-114-95-PROPAGATION (anchor literal across canonical fleet) + CM-WORKABLE-ITEMS-DB-TRACKED (git ls-files docs/workable_items.db returns non-empty if DB exists on disk + .gitignore does NOT exclude it). Paired §1.1 meta-test mutation adds the DB to .gitignore → gate FAILs.

Composes with §11.4.93 (SQLite-SSoT — §11.4.95 amends its tracking-policy clause), §11.4.30 (.gitignore discipline — explicit carve-out anchored here), §11.4.77 (regeneration mechanism — does NOT apply to the DB; the DB IS the source), §2.1 (multi-upstream push — DB pushed to every mirror), §9.2 (data safety — destructive DB ops require hardlinked-backup + operator authorization), §11.4.41 (force-push merge-first — applies if DB history ever needs rewrite).

Canonical authority: this Constitution.md §11.4.95 in the HelixConstitution submodule.

Non-compliance is a release blocker. A consuming project that gitignores its workable-items DB is severity-equivalent to a §11.4 PASS-bluff at the source-of-truth layer — it implies SSoT exists when it is in fact ephemeral local-only state.

§11.4.96 — Safe-parallel-work-with-long-build catalogue + mandate (User mandate, 2026-05-27)

Forensic anchor — verbatim user mandate (2026-05-27):

“Are there except AOSP build process any other active jobs being done at the moment? Can we work on something in parallel while build is in progress so we slowly cleanup our slate? Anything workable which will not affect the build process or final image we generate for the flashin. If yes, and that was not already clear by our root Constitution, we should add all mandatory details into it ... Make sure we do as much as possible work in background in parallel with main work stream and oreferably using subagents-driven approach!”

§11.4.94 mandates “always check parallel-work feasibility before idle”; §11.4.89 mandates background-test discipline; §11.4.58 defines the parallel-work-unit pipeline. §11.4.96 is the **operational catalogue** that enumerates, for the canonical long-running workload (5-7 h AOSP containerised build per §12.9), exactly what CAN safely run in parallel and what MUST NOT — closing the ambiguity gap §11.4.94 left implicit.

Closed-set categories — SAFE during AOSP containerised build:

- **(A) Markdown / documentation work** anywhere under docs/, *.md, README, CLAUDE.md, AGENTS.md, QWEN.md, Status.md / Status_Summary.md, Issues.md / Fixed.md / their Summary siblings, changelogs, research notes. The build container mounts source read-only at a snapshot moment + reads only the build-relevant files (Android.bp / Android.mk / source code / vendor binaries). MD edits do NOT propagate into the build’s read scope.
- **(B) Generator + helper script work** under scripts/, scripts/testing/, scripts/llm/, scripts/firebase/, etc. — these are host-side helpers NEVER consumed by the AOSP build. Safe to edit + commit + push freely.
- **(C) Pre-build / meta-test gate authoring + paired §1.1 mutations** under device/rockchip/rk3588/tests/pre_build_verification.sh + scripts/testing/meta_test_false_positive_proof.sh — these run host-side via bash pre_build_verification.sh ; the AOSP build itself does NOT execute them inline. Safe to edit during build; the NEXT pre-build run (post-current-build) picks up the new gates.
- **(D) On-device tests** under device/rockchip/rk3588/tests/test_*.sh — these run on D3/D4 via ADB post-flash. Authoring new test scripts during build does NOT affect the build (the build packages the test corpus from the source tree but does so once at copy-out time; scripts added after the copy-out are absent from THIS image but present in source for the NEXT image).
- **(E) Constitution submodule edits + push** — separate git workspace at constitution/, not on the AOSP build’s path. Edits + push to all 6 constitution remotes safe at any time.
- **(F) Project submodule commit + push** to either the parent or owned submodules’ remotes — per §11.4.88 background push, this is OS-level orthogonal to the container.
- **(G) D3 / D4 live-ADB probes** (read-only dumsys / getprop / cat /proc/asound/... / screencap / logcat -d) — these query device state without changing it; do NOT affect the host build.
- **(H) Subagent dispatch** per §11.4.20 / §11.4.70 — each subagent runs in isolated context window; multiple subagents in same checkout coordinate via §11.4.84 working-tree quiescence (lockfile or git worktree per disjoint scope). Read-only analysis subagents (obsolescence audit, summary clarity sweep, code-review) are uniformly safe.
- **(I) Web research + external API queries** (with §11.4.10 credential discipline preserved) — operator-relevant browsing, ASUS QVL lookups, vendor spec verification, package-registry version checks.
- **(J) Workable-items DB operations** per §11.4.93 + §11.4.95 — workable-items sync (when implemented), workable-items validate, workable-items report — read/write SQLite at docs/workable_items.db, orthogonal to build.
- **(K) Pre-build verification + meta-test execution** — RUN per §11.4.89 background-test discipline. The build container reads source; pre-build runs host-side reading the same source. No conflict at filesystem layer (read-only intersection).

Closed-set — UNSAFE during AOSP containerised build:

- **(α) git checkout / git reset --hard / git clean -df** on the source tree — would change files the build is reading; introduces inconsistency. Defer until build completes OR use a git worktree per §11.4.84.
- **(β) Mass file deletions or renames** under device/, frameworks/, hardware/, vendor/, kernel-5.10/, external/, packages/, bionic/, bootable/, build/, system/, cts/, tools/, prebuilts/, etc. — anything the AOSP build potentially reads. Apply ONLY after build completes.
- **(γ) Submodule pointer updates that affect built APKs** — bumping device/rockchip/atmosphere/presenter or smarttube-player or vlc-player etc. pointer changes the binary content the build packages. Defer pointer bumps to between-build windows OR plan for the NEXT build to pick them up.
- **(δ) out/ directory mutations** — the build's output directory. Never touch concurrently.
- **(ε) make clean / m clobber / rm -rf out/** while build runs — direct conflict; will likely crash the build.
- **(ζ) Container destruction** (podman stop atmosphere-aosp-build, podman pod rm) — terminates the build mid-flight. Operator-explicit only.
- **(η) Disk-filling operations** (large file downloads, container image builds, NVMe duplication) that push the host below the §12.9 free-space minimum — could OOM the build's writable layer.
- **(θ) Host-session-safety breaches** per §12 — suspend / hibernate / logout / unbounded memory ops in user.slice — kill the build's parent cgroup.

Conductor responsibility (per §11.4.94 already; §11.4.96 makes the catalogue explicit):

- Before EVERY pause point during a long build, the conductor MUST consult this catalogue, identify all (A)-(K) items in the priority queue per §11.4.42 + §11.4.72, dispatch at least one per §11.4.20 / §11.4.70 subagent-driven default + §11.4.89 background-task discipline. “Build is running, nothing else to do” is NEVER true per §11.4.94 + this catalogue.

Subagent-driven default for catalogue items:

- 1. MD work + (E) constitution edits + (F) commits: conductor-direct (lightweight + critical-state sequencing).
- 1. generator/helper work + (C) gate authoring + (D) on-device test authoring + (G) live-ADB probes + (I) web research + (J) DB ops: subagent-dispatchable per §11.4.20.
- 1. is itself the subagent dispatch — composes with the rest.
- 1. pre-build / meta-test execution: backgrounded per §11.4.89, conductor polls outcome.

Pre-build gate CM-COVENANT-114-96-PROPAGATION enforces this anchor literal across the canonical fleet. Pre-build gate CM-PARALLEL-WORK-DURING-BUILD-AUDIT (when implemented) audits recent commits during AOSP build windows for the parallel-work output (commits landing during a build window with timestamps inside qa-results/aosp_build_*.log start/end). Paired §1.1 meta-test mutations strip the catalogue literals → gates FAIL.

Composes with §11.4.20 (subagent-driven) + §11.4.42 (iteration priority) + §11.4.58 (parallel PWU) + §11.4.70 (subagent default) + §11.4.72 (audio top-priority) + §11.4.82 (iteration-speedup) + §11.4.84 (working-tree quiescence —

esp. for subagent parallelism) + §11.4.85 (stress + chaos can run in parallel) + §11.4.87 (endless-loop zero-idle) + §11.4.88 (background-push) + §11.4.89 (background-test) + §11.4.92 (multi-pass evaluation per parallel branch) + §11.4.93 / §11.4.95 (workable-items DB) + §11.4.94 (parallel-work survey mandate) + §12.6 / §12.7 / §12.8 / §12.9 (host-session safety + build cap).

Canonical authority: this Constitution.md §11.4.96 in the HelixConstitution submodule.

Non-compliance is a release blocker. Conductor scheduling a wake/sleep during an active AOSP build without first dispatching at least one (A)-(K) catalogue item from the priority queue is a §11.4.94 + §11.4.96 violation = §11.4 PASS-bluff at the operating-mode layer.

§11.4.97 — Maximum-use-of-idle-time mandate + progress-update cadence (User mandate, 2026-05-27)

Forensic anchor — verbatim user mandate (2026-05-27):

“keep it working, we should do as much as possible, if not it all but as much as we can as long as there is iddle time! it MUST be used! make sure that it is clear in constitution! keep us updated about all progress and all phisycal proofs and gathered data as you progress through all open workable items!”

§11.4.97 strengthens §11.4.87 (endless-loop) + §11.4.94 (zero-idle priority-first) + §11.4.96 (safe-parallel-with-build catalogue) with an **explicit maximum-use-of-idle-time mandate** AND a **progress-update cadence mandate**:

(A) Maximum-use idle-time mandate. Every minute of conductor idle time during which (i) the conductor has work it could autonomously progress AND (ii) is NOT genuinely blocked on external dependency (operator hardware action, network upstream, host-session safety, build/test completion) is a §11.4.97 violation. “As much as possible, if not it all but as much as we can” is the operative phrasing — the conductor MUST dispatch work continuously through the entire idle window, not just sample-check at scheduled wakes.

(B) Progress-update cadence mandate. The conductor MUST emit operator-facing progress updates at the following natural milestone boundaries (no operator prompt required):

- Every commit that lands (HEAD advance) — 1-line “what just landed”.
- Every subagent return — 1-line “subagent X returned, integrated”.
- Every constitutional anchor landed — 1-line “§X.Y.Z propagated to 6 remotes”.
- Every gathered physical proof — 1-line “captured evidence at qa-results”.
- Every milestone closure (Issues→Fixed migration, Obsolete classification, etc.) — 1-line.

Progress updates are CONCISE (1-3 lines each); the operator’s mental model stays current without conductor overhead.

(C) Continuous physical-proof gathering mandate. Per §11.4.5 + §11.4.6 + §11.4.69 — every autonomous-fixable item the conductor closes MUST gather positive captured-evidence: ADB probes, dumpsys captures, screenshots, ALSA

hw_params, sysfs reads, file-content hashes, runtime metric snapshots. The evidence is committed alongside the closure narrative. Per §11.4.93 SQLite-SSoT, the evidence path goes into the item_history.evidence_path column when DB integration lands.

(D) Per-anchor composition. §11.4.97 binds together: §11.4.5 (captured-evidence quality), §11.4.6 (no-guessing — every progress claim cites evidence), §11.4.13 (sink-side evidence when applicable), §11.4.20 (subagent-driven default for parallel scope), §11.4.27 (no-fakes-beyond-unit — closures cited with real captured evidence), §11.4.42 (priority queue), §11.4.50 (deterministic consistency — N=3 where applicable), §11.4.52 (autonomous validation), §11.4.69 (sink-side taxonomy), §11.4.70 (subagent-driven), §11.4.72 (audio top-priority), §11.4.83 (docs/qa transcript), §11.4.85 (stress + chaos), §11.4.87 (endless-loop), §11.4.88 (background-push), §11.4.89 (background-test), §11.4.94 (zero-idle), §11.4.96 (safe-parallel catalogue).

(E) Idle-only-when-genuinely-blocked. Same closed-set as §11.4.94(A): operator STOP, external dependency the conductor cannot accelerate, host-session-safety demand. Per §11.4.96 the safe-during-build catalogue (A-K) covers what's progressable; the conductor MUST exhaust the catalogue before scheduling sleep.

Operator-facing visibility. §11.4.97's progress-update cadence is the contract by which the operator stays informed without prompting. The conductor does NOT wait for operator status checks — milestone updates emit autonomously.

Pre-build gate CM-COVENANT-114-97-PROPAGATION enforces this anchor literal across the canonical fleet. Pre-build gate CM-IDLE-TIME-AUDIT (when implemented) audits the conductor's wake/sleep schedule across recent sessions for unjustified idle periods. Paired §1.1 meta-test mutations strip the load-bearing literals → gates FAIL.

Composes with every anchor in §11.4 (this is the operating-mode capstone): §11.4.5 + §11.4.6 + §11.4.13 + §11.4.20 + §11.4.27 + §11.4.42 + §11.4.50 + §11.4.52 + §11.4.69 + §11.4.70 + §11.4.72 + §11.4.83 + §11.4.85 + §11.4.87 + §11.4.88 + §11.4.89 + §11.4.94 + §11.4.96.

Canonical authority: this Constitution.md §11.4.97 in the HelixConstitution submodule.

Non-compliance is a release blocker. Conductor idle when progressable work remains, OR conductor failing to emit milestone progress updates, is severity-equivalent to a §11.4 PASS-bluff at the operating-mode layer.

§11.4.98 — Full-Automation Anti-Bluff Mandate — Live tests MUST be re-runnable end-to-end without manual intervention (User mandate, 2026-05-28)

Forensic anchor — verbatim user mandate (2026-05-28):

“Make sure we have full automation testing of all scenarios with real bot, main group and users without any manual intervention or contribution of real user! Everything MUST BE fully automatic and autonomous! These tests MUST BE able to rerun endless times when needed! This is important to be

done like this! It is critical! Continue all work and make this happen! We need such full automation testing so the whole System MUST BE fully validated and verified before it is integrated to our main projects! Make sure there is no false positives in testing! Every test and its results MUST obtain real proofs of everything working! No bluff is allowed! IMPORTANT: Make sure that all existing tests and Challenges do work in anti-bluff manner - they MUST confirm that all tested codebase really works as expected!”

§11.4.98 composes with §11.4 / §11.4.2 / §11.4.5 / §11.4.50 / §11.4.85 / §11.4.87 / §11.4.89 / §11.4.94 — closes the **manual-intervention gap** that they did not explicitly forbid. A live/integration/e2e/Challenge test which requires a human action during execution (typing a chat message, clicking a UI, hand-triggering a webhook, manually attaching a file, anything beyond the test starting up and reporting PASS/FAIL on its own) is **by definition a §11.4 PASS-bluff at the automation layer**, regardless of how thorough the manual run is. The reason: such a test cannot run continuously in CI, cannot validate regressions between manual runs, and the human dependency masks any drift between code changes and what the test exercises.

(A) Binding rule. Every test that this Constitution governs — unit / integration / e2e / Challenge / stress / chaos / live — MUST be fully self-driving end-to-end. The test process must accept inputs from configuration alone, exercise the load-bearing code path autonomously, and report PASS/FAIL/SKIP-with-reason without any further human action.

(B) Single permissible exception — one-time credential bootstrap. Configuration performed OUTSIDE test execution is acceptable: populating `.env` from a vault, exporting shell env vars in `~/.bashrc`, OAuth approval at first install, MTPROTO session activation at first run. This is configuration, not test driving. Once credentials are present, the test MUST run end-to-end without any further human action. The configuration step is itself a §11.4.10 credential-handling concern, not a §11.4.98 automation concern.

(C) Concrete requirements for live messenger / channel / agent tests (the surface that exposed this mandate):

1. **No “operator MUST type a message” prompts in test scripts.** Tests requiring inbound messages MUST drive those messages programmatically — via a separate user account (MTPROTO for Telegram, real-user-token API for Slack, IMAP-test-account for email, etc.), via a webhook fixture, or via an in-process loopback. Never via human keystrokes during test execution.
2. **No hard-coded session UUIDs that collide with the active dev session.** Test runs MUST use a dedicated test-only session that the production / dev session is NOT simultaneously using. (Lesson from Herald 2026-05-28: `claude --resume <UUID>` invoked against the same session ID the dev session is using returns exit -1 with no output — a silent collision that fails for non-obvious reasons.)
3. **No 60-second human-response windows.** These create flake/timing races that are themselves §11.4.50 determinism violations: a single test invocation’s PASS/FAIL depends on a human reaction-time variable, not on the code under test.
4. **Re-runability proof.** Every live test MUST PASS at least 3 consecutive automated invocations (`-count=3` for Go, `repeat 3` equivalent for other runtimes) with no state cleanup between runs. Persistent side effects (chats,

files, DB rows, attached objects, queued events) MUST be cleaned by the test itself in a defer/teardown block.

5. **§11.4.98 obsolescence audit.** Every existing test that ships under this Constitution MUST be classified as either “self-driving” (COMPLIANT) or “manual-dependency” (NON-COMPLIANT — must be rewritten or marked §11.4.90 Obsolete). The audit is a release-gate item — a release that ships a NON-COMPLIANT test still pretending to PASS is blocked.
6. **No false-positive PASS.** A test PASS that does NOT exercise the load-bearing code path — e.g. a “live” test that skips silently because env vars are missing yet reports PASS, or a “live” test where the captured evidence is from a prior cached run — is itself a §11.4 PASS-bluff at the gating layer. SKIP-with-reason is the §11.4.3 correct posture; SKIP-reported-as-PASS or stale-evidence-reported-as-fresh is forbidden.

(D) Composition. §11.4.85 (stress + chaos) + §11.4.89 (background-test execution) + §11.4.87 (endless-loop autonomous work) + §11.4.94 (zero-idle parallel-by-default) + §11.4.98 (full-automation) together make this Constitution’s testing surface a continuously-validated, fully-automated, non-flake, anti-bluff regime. Each closes a different gap; remove any one and the whole property collapses. §11.4.98 specifically closes the manual-intervention gap: §11.4 + §11.4.85 + §11.4.89 + §11.4.87 + §11.4.94 do not forbid a test from requiring human action, only from skipping the load-bearing path. §11.4.98 makes the human dependency itself a §11.4 violation.

(E) Inheritance per §11.4.35. §11.4.98 propagates to every consuming repository via the existing inheritance gate: each project’s CLAUDE.md / AGENTS.md / QWEN.md MUST carry a short-form restatement citing the literal anchor 11.4.98. The pre-build gate CM-COVENANT-114-98-PROPAGATION (when implemented) enforces this literal anchor presence across the canonical fleet. Paired §1.1 meta-test mutations strip the load-bearing literal → gates FAIL.

(F) Enforcement. A commit that adds or modifies a test that requires manual human action during execution is blocked at release-gate. A test classified as “manual-dependency” in the §11.4.98 audit that has not been rewritten within 30 days of classification graduates to §11.4.90 Obsolete and is removed from the active test suite (not deleted — preserved with Obsolete-Details: per §11.4.90, citing §11.4.98 as the obsolescence reason).

Canonical authority: this Constitution.md §11.4.98 in the HelixConstitution submodule. All consuming projects (Herald, ATMOSphere, future) restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker. No --manual-test-OK, --skip-114-98-audit, --bluff-tolerance-temporary flag exists. The 2026-05-28 user mandate is unambiguous: “No bluff is allowed!”

§11.4.99 — Latest-Source Documentation Cross-Reference Mandate — instructions, guides, and manuals MUST be verified against the latest official online sources BEFORE publication (User mandate, 2026-05-28)

Forensic anchor — verbatim user mandate (2026-05-28):

“Make sure we ALWAYS check against latest versions of services we use web / online docs before creating instructions! This situation is illustration of how we can misguide ourselves or get banned! Add this all important generic / general points as proper mandatory rules when we are creating documentation or guides! These are mandatory rules / constraints and the result is consistency and safety of created instructions, guides and manuals!”

Case study that anchored this rule (Herald 2026-05-28). The first-draft Herald MTProto setup guide (committed at 35fc10c / fb00354 / f089dd6) recommended VoIP / Google Voice / Twilio / TextNow numbers as a budget-friendly fallback option AND omitted the critical recover@telegram.org pre-login email step. Both pieces of guidance directly contradicted (a) Telegram’s official documentation at https://core.telegram.org/api/obtaining_api_id (“all accounts that sign up or log in using unofficial Telegram clients are automatically put under observation”) AND (b) the gotd/td maintainer’s “How to not get banned?” guide (vendored at submodules/gotd-td/.github/SUPPORT.md). An operator who followed the original guide could have had their Telegram account permanently banned with no appeal — the very harm the documentation was supposed to help them avoid. The corrected guide landed at 8470ba7 after a forced cross-reference against the latest official sources. This case study is forensic evidence that misguidance-by-stale-docs is the same severity class as a §11.4 PASS-bluff at the documentation layer.

Composition. This anchor composes with §11.4.4 (test-interrupt-on-discovery), §11.4.5 (captured-evidence), §11.4.8 (deep-research validation), §11.4.92 (multi-pass change-evaluation Pass 4), §107 (anti-bluff), and §11.4.98 (full-automation). It closes a specific gap §11.4.92 Pass 4 alludes to but does not explicitly mandate: **every operator-facing instruction, guide, manual, troubleshooting cookbook, or setup walkthrough MUST be cross-referenced against the LATEST official online documentation of the service / library being documented BEFORE the document is committed.**

(A) Binding rule — pre-commit cross-reference. Before committing any document that contains operator-actionable steps (setup walkthroughs, integration guides, API how-tos, credential acquisition, security configurations, troubleshooting cookbooks, billing-policy documentation, terms-of-service summaries, any external-service interaction guide), the author MUST:

1. **Fetch the latest official online documentation** of the third-party service / library being documented via WebFetch, direct browsing, the service’s own MCP server (if available), or equivalent authoritative real-time source. Do NOT rely on training data, memory, prior assumptions, or older committed docs as the source of truth.
2. **Cross-reference each instruction in the new document** against that source. For each step the operator will take, verify: (a) the service still supports that action; (b) the form fields / parameters / endpoints are still the same; (c) any new constraints, deprecations, or warnings the service published since the doc’s last verification.
3. **Seek secondary authoritative sources when the official documentation is sparse / silent** on a critical requirement. Examples: the library maintainer’s official guides (submodules/<lib>/README.md, SUPPORT.md, SECURITY.md), the service’s official blog / changelog, the service’s official support email / channel responses, community-vetted FAQs (Stack Overflow accepted answers, official Discord pinned messages).
4. **Cite the source URL + the date checked at the bottom of the document** in a ## Sources verified section. Example: Sources verified 2026-05-28:

https://core.telegram.org/api/obtaining_api_id + submodules/gotd-td/.github/SUPPORT.md. The citation is the audit trail; a doc without it is by §11.4.99 definition NOT verified.

5. **Cite the cross-reference in the commit message footer.** A commit that adds or modifies operator-facing instructions MUST include a Sources verified <date>: <urls> footer line so the verification trail is reachable from git log queries.

(B) Negative-finding documentation is required. If the cross-reference reveals the official source is silent / contradictory / outdated on a question the new document needs to answer, the document MUST explicitly note that gap so the next reader does not assume the absence of contradiction means authoritative agreement. Example phrasing the Herald MTPROTO guide uses: *“The official Telegram docs do not enumerate App title validation rules; the constraints below were verified against the gotd/td community channel + empirical operator testing 2026-05-28.”*

(C) Re-verification cadence. Documents older than 6 months without re-verification are STALE — operators MUST NOT trust them for fresh action without re-running the cross-reference. Documents MUST be re-verified:

1. **Before being cited as the authority for an operator-action campaign** (the operator’s command to “follow this guide” is the trigger).
2. **At every major release boundary** (vN.0.0 tag — the release-gate sweep includes a §11.4.99 freshness audit).
3. **Whenever the documented service publishes a breaking-change announcement** (the agent that knows about the announcement MUST queue the re-verification).
4. **When an operator reports an error following the guide** (the case-study anchoring is automatic and triggers immediate re-verification).

(D) Service-specific risk-classifications. Documentation for the following service families MUST include explicit safety warnings cross-referenced against the latest published policies, with a §11.4.99 “Sources verified” date NEVER older than 90 days:

Family	Specific risks the documentation MUST address
Telegram / WhatsApp / messenger APIs (unofficial clients)	Anti-abuse-system observation; one-phone-one-app-id limits; ban-on-VoIP policies; rate-limit floods; user-impersonation risks; pre-login declaration emails.
Cloud provider APIs (AWS, GCP, Azure, Cloudflare)	Billing blast-radius; account-lock policies; service-quota limits; least-privilege IAM; data-residency obligations.
Payment systems (Stripe, PayPal, banks, Mercado Pago)	KYC/AML compliance; PCI-DSS data-handling rules; webhook signature verification; idempotency keys; refund/chargeback flows.
AI / LLM providers (Anthropic, OpenAI, Google AI)	Terms-of-service violation triggers; rate-limit + quota policies; data-

Family	Specific risks the documentation MUST address
	retention defaults; safety-classifier outcomes; PII-in-prompt risks.
Code-hosting services (GitHub, GitLab, GitFlic, GitVerse, Bitbucket)	Token-leak revocation policies; force-push to protected-branch policies; rate-limit-on-API-tokens; secret-scanning bot behaviour.
OS / package managers (apt, brew, npm, pip, cargo)	Supply-chain compromise vectors; signature verification; lockfile discipline; mirror-trust policies.

This list is **NOT exhaustive** — when a new external service is documented, the §11.4.99 author judges whether the service has comparable risk surface; if yes, the same safety-warning requirement applies.

(E) Composition with §11.4.92 multi-pass change-evaluation. §11.4.92 Pass 4 (deep-research validation) already requires “external precedent or literal ‘NO external solution found’”. §11.4.99 strengthens this specifically for the operator-facing-instructions class: deep research is mandatory for ALL such docs, NOT just non-trivial code changes. The agent that authors an instruction guide CANNOT cite §11.4.92 Pass 4 as a substitute for §11.4.99 — the two pass independently.

(F) Inheritance per §11.4.35. Every consuming repository’s CLAUDE.md / AGENTS.md / QWEN.md MUST carry a short-form restatement citing the literal anchor 11.4.99. The pre-build gate CM-COVENANT-114-99-PROPAGATION (when implemented) enforces this literal anchor presence across the canonical fleet. Paired §1.1 meta-test mutations strip the load-bearing literal → gates FAIL.

(G) Enforcement. A commit that adds or modifies operator-facing instruction documentation without (a) a “Sources verified ” footer in the document itself AND (b) a Sources verified line in the commit-message footer is blocked at release-gate. A document with operator-actionable steps that becomes stale (older than 6 months without re-verification, or 90 days for risk-classified services per §(D)) graduates to §11.4.90 Obsolete with Obsolete-Details: Reason=stale-documentation; superseded-by=<replacement-doc-or-rewrite> after the 30-day grace window per §11.4.90.

Canonical authority: this Constitution.md §11.4.99 in the HelixConstitution submodule. All consuming projects (Herald, ATMOSphere, future) restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker. No --skip-source-check / --documentation-freshness-optional / --cite-sources-later / --trust-prior-doc-as-authoritative flag exists. The 2026-05-28 user mandate is unambiguous: the result MUST be “consistency and safety of created instructions, guides and manuals”.

§11.4.100 — RETIRED. Demoted to consumer project (ATMOSphere video-color/visual-quality fidelity) per §11.4.17/§11.4.35 — project-specific (RK3588/MPV/Arvus), not universal. See the consuming project’s Constitution/CLAUDE/AGENTS/QWEN.

§11.4.101 — Autonomous-decision-over-blocking mandate (User mandate, 2026-05-28)

Short tag: autonomous-decision-over-blocking.

Forensic anchor — verbatim user mandate (2026-05-28):

“when working in endless working loop fully autonomously try to decide most properly about points which would block execution and wait for us. If we haven’t answered now work would be blocked whole night! If possible and if that will not cause any issues make proper and most reliable and safe decision so we achieve maximal efficiency and work gets fully done!”

The mandate. When operating in an autonomous / endless-loop mode (per §11.4.87), the agent **MUST** minimize operator-blocking and instead make the safe, reliable, reversible decision itself — so work is **NOT** stalled (e.g. overnight) waiting for input. The §11.4.87 endless-loop covenant told the agent to keep working; §11.4.101 tells it **HOW** to handle the decision points that would otherwise force it to stop and wait. A wrong bias toward blocking-when-it-was-safe-to-decide wastes operator time (the forensic harm: “blocked whole night”); a wrong bias toward deciding-when-it-should-have-blocked risks irreversible harm. The closed-set decision rule below is the precise boundary between the two.

Decision rule (closed-set — the agent MAY proceed autonomously when ALL hold):

1. the action is **reversible** OR has a captured pre-op backup per §9.2 (hardlinked .git mirror, recorded refs/tags/submodule pointers);
2. the agent can **determine the safe choice from captured evidence** per §11.4.6 (no guessing — **LIKELY** / probably / seems is **NOT** a safe-choice determination; the evidence must state the correct choice as fact, or the condition is not met);
3. a wrong choice’s **blast radius is bounded AND recoverable** — the damage is local, undoable, and does not cascade beyond the work unit;
4. it **composes with the existing covenants** — anti-bluff §11.4 (no PASS-bluff to skip the decision), host-session-safety §12 (no decision that suspends / logs-out / overruns the §12.6 memory budget), data-safety §9 (no destructive op without the §9 protocol).

Block-only-when rule (the agent MUST BLOCK — ask via the §11.4.66 interactive mechanism — ONLY when ALL of the following hold): the action is **irreversible AND high-blast-radius AND the safe choice cannot be determined from evidence**. Canonical block-required examples: external-account state the agent cannot inspect (registration, certification, billing portals), hardware the agent cannot physically access, destructive repository ops without a backup, force-push (which independently requires §9.2 + §11.4.41 authorisation), spending money or sending messages / data to third parties. Operator-blocked per §11.4.21 remains the last-resort status and is reached only after this rule fires **AND** the §11.4.21 self-resolution-exhaustion audit is complete.

Maximize-progress-while-blocked rule. When a block **IS** unavoidable, the agent **MUST** still maximize progress on every **NON**-blocked item in parallel per §11.4.87 + §11.4.94 (zero-idle priority-first) rather than idling — the blocked decision parks

one work unit, it does not pause the loop. Posing the §11.4.66 question and continuing other work is the correct behaviour; posing the question and going idle is a §11.4.94 + §11.4.97 violation.

Composes with §11.4.6 (no-guessing — sub-rule (b) IS the no-guessing test applied to the decision itself), §11.4.21 (operator-blocked is the last-resort status the block-rule routes into, after self-resolution exhaustion), §11.4.40 (a release-tag decision is irreversible-and-high-blast-radius → always blocks per the block-rule), §11.4.41 (force-push is an explicit block-required example), §11.4.66 (the interactive-clarification mechanism the block-rule uses), §11.4.87 (the endless-loop covenant §11.4.101 refines — §11.4.87 says keep working, §11.4.101 says how to clear the decision points), §11.4.94 (zero-idle — maximize-progress-while-blocked), §9.2 (backup is the reversibility substitute in sub-rule (a)), §12 (host-safety bounds every autonomous decision).

Classification: universal (§11.4.17) — autonomous-decision-vs-block boundary is a reusable discipline for ANY project running an agent in an endless / unattended loop, hardware- and domain-agnostic.

4-layer coverage per §11.4.4(b). Propagation gate CM-COVENANT-114-101-PROPAGATION enforces the literal anchor 11.4.101 across the canonical consumer fleet (parent + owned-submodule CLAUDE.md / AGENTS.md / QWEN.md). Paired §1.1 meta-test mutation strips the 11.4.101 literal from a consumer file → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md §11.4.101 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no --always-block-on-decision, --never-decide-autonomously, --skip-decision-rule, --block-without-self-resolution flag exists. The 2026-05-28 user mandate is unambiguous: in autonomous mode the agent makes the safe reversible decision and keeps working — it does NOT block the operator overnight on a choice it had the evidence to make safely.

§11.4.102 — Mandatory systematic-debugging activation + always-loaded skill-discovery + plugin-dependency availability (User mandate, 2026-05-29)

Short tag: systematic-debugging-always-on.

Forensic anchor — verbatim user mandate (2026-05-29):

“Make sure that we ALWAYS trigger / start the”/superpowers:systematic-debugging” skills when any issues happen! If this is possible to activate and use in this situations out of the box when we spot problems / issues / bugs / misalignments / inconsistencies we MUST activate the skill(s) and make strongest efforts in full in depth analysis / debugging and determine root causes of all problem or obtain relevant data and information we need! ... we MUST make sure that “/using-superpowers” skill is ALWAYS loaded,

applied and used! All dependencies (plugins) that Claude Code or other market places are offering MUST BE installed if these are not already available for loading and use!”

This anchor binds three cooperating invariants. Together they make the difference between an agent that **guesses-and-retries** (symptom-patching, “transient”, “flaky”, “probably a race”) and an agent that **investigates-to-root-cause-first** mechanical, every single time a problem surfaces. The §11.4.6 no-guessing mandate said *do not state an unproven cause*; §11.4.102 says *the moment a problem appears, the very first action is to launch the structured root-cause investigation skill — not to propose a fix*.

(A) Mandatory systematic-debugging activation on ANY spotted problem. On ANY spotted issue, bug, test failure, gate failure, regression, misalignment, inconsistency, unexpected behaviour, or contradiction-between-evidence-and-expectation, the agent MUST activate the structured root-cause-investigation skill — `superpowers:systematic-debugging` on Claude Code, or the platform-equivalent structured-debugging discipline on any other runtime — **BEFORE proposing, writing, or applying any fix**. This is the **Iron Law: NO FIXES WITHOUT ROOT CAUSE INVESTIGATION FIRST**. The investigation MUST run its full four-phase arc: (1) **root-cause phase** — read the actual error / failing assertion / divergent evidence, reproduce deterministically, gather the facts the failure exposes; (2) **pattern phase** — identify what class of defect this is, search for the same pattern elsewhere in the codebase; (3) **hypothesis phase** — form a falsifiable hypothesis about the cause, then prove or disprove it with captured evidence (never **LIKELY** / probably / seems per §11.4.6); (4) **implementation phase** — only after the cause is a **FACT**, design the fix against the proven root cause. Guess-and-retry, symptom-patching, “let me just try changing this and see”, and re-running a failed test hoping it passes (“it’s probably transient / flaky”) **WITHOUT** a completed root-cause investigation are themselves §11.4.102 violations. Real-world signal this anchor closes: declaring a failure transient / flaky / intermittent / probably a timing issue **without captured forensic evidence proving that classification** is simultaneously a §11.4.6 (no-guessing) violation AND a §11.4.7 (demotion-evidence) violation — §11.4.102 makes the corrective response mechanical: the classification is forbidden until the systematic-debugging arc has captured the evidence that proves it.

(B) Mandatory always-loaded skill-discovery discipline (using-superpowers). The skill-discovery meta-skill — `superpowers:using-superpowers` on Claude Code, or the platform-equivalent skill-index / capability-discovery discipline on any other runtime — MUST be loaded and applied at session start and consulted before any task is begun. Its operative rule: **before acting on ANY request, survey the available skills/capabilities; if ANY skill could apply to the task at hand — even at a 1% chance of relevance — that skill MUST be invoked rather than the agent improvising from memory**. This guarantees the agent reaches for the purpose-built discipline (systematic-debugging for a bug, the brainstorming skill before creative work, the test-driven-development skill before writing a fix, etc.) instead of ad-hoc behaviour. Skipping skill-discovery at session start, or improvising a task that a loaded skill was built to handle, is a §11.4.102 violation.

(C) Mandatory plugin / dependency availability. Every skill plugin, marketplace package, or capability dependency that the project relies on — offered by Claude Code, its skill marketplaces, or the equivalent plugin/extension ecosystem of any other runtime the project uses — MUST be installed and loadable on the host **BEFORE** the dependent work proceeds. A missing plugin/dependency that blocks a

mandated skill (e.g. the superpowers plugin absent so systematic-debugging cannot launch) is a **release-blocker** until it is installed and confirmed loadable. Install mechanism (generic; runtime supplies the concrete command): the runtime's own plugin/marketplace install path — for Claude Code, the in-session /plugin marketplace flow (add the marketplace, install the plugin, confirm the skill appears in the available-skills list); for other runtimes, the documented package/extension installer (npm/pip/cargo/the editor's extension manager). The availability check is itself anti-bluff: a plugin claimed “installed” without the dependent skill actually appearing in the loaded-skills/capabilities list is a §11.4.6 (no-guessing) violation — confirm by observing the skill in the live capability list, not by assuming the install succeeded (install-command exit 0 ≠ skill loadable, mirroring the §11.4.80 npm exit 0 ≠ working binary lesson).

Composes with §11.4.4 (test-interrupt-on-discovery — §11.4.4 says STOP the cycle the moment a defect is discovered; §11.4.102 says the FIRST action after stopping is to launch systematic-debugging), §11.4.6 (no-guessing — §11.4.102(A) is the procedural enforcement that produces the facts §11.4.6 demands; the forbidden LIKELY / transient / flaky vocabulary is exactly what guess-and-retry emits), §11.4.7 (demotion-evidence — a FAIL→lower-severity demotion requires same-conditions evidence, which is precisely what the systematic-debugging arc captures), §11.4.8 (deep-web-research — runs inside the hypothesis phase), §11.4.43 (TDD-fix — the RED test is written against the root cause the investigation proved), §11.4.70 (subagent-driven — a systematic-debugging investigation is a natural subagent dispatch), §11.4.82(A) (iteration-speedup — already mandates Phase-1 forensic before any speculative patch; §11.4.102 generalises that from “source patch” to ANY fix and adds the always-loaded skill-discovery + plugin-availability invariants), §11.4.92 (multi-pass change-evaluation — the investigation feeds Pass 1 + Pass 4).

Classification: universal (§11.4.17) — root-cause-before-fix, skill-discovery-before-task, and dependency-availability are reusable disciplines for ANY project on ANY runtime that exposes a skill/plugin ecosystem; the concrete skill names (superpowers:systematic-debugging, superpowers:using-superpowers) are the Claude Code instantiation, with the generic “or platform-equivalent” wording carrying the rule to other runtimes per §11.4.81 cross-platform-parity reasoning.

4-layer coverage per §11.4.4(b). Propagation gate CM-COVENANT-114-102-PROPAGATION enforces the literal anchor 11.4.102 across the canonical consumer fleet (parent + owned-submodule CLAUDE.md / AGENTS.md / QWEN.md). Paired §1.1 meta-test mutation strips the 11.4.102 literal from a consumer file → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md §11.4.102 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no --skip-systematic-debugging, --guess-and-retry-OK, --symptom-patch-permitted, --skip-skill-discovery, --plugin-optional, --missing-plugin-is-warning flag exists. The 2026-05-29 user mandate is unambiguous: the moment a problem is spotted the structured root-cause investigation skill activates, using-superpowers is always loaded, and every mandated plugin/dependency is installed and loadable — there is no path by which a fix is proposed before the cause is proven.

§11.4.103 — Continuous parallel-stream working routine (User mandate, 2026-05-29)

Short tag: continuous-parallel-stream-routine.

Forensic anchor — verbatim user mandate (2026-05-29):

“Do this working approach continuously and make it part of regular working routine and add it to the Constitution Submodule documented fully.”

This anchor promotes the proven multi-stream operating pattern — repeatedly demonstrated effective across prior autonomous rounds — from an ad-hoc practice into the project’s **standing default working routine**. §11.4.87 told the agent to keep working in an endless loop; §11.4.94 told it never to sit idle when something is progressable; §11.4.88 / §11.4.89 detached push and test execution; §11.4.96 catalogued what is safe to do during a long build. §11.4.103 binds those into a single always-on routine and adds the load-bearing invariant the prior anchors only implied: **the main work stream MUST always stay FREE**. It is the conductor’s standing operating mode, not a per-request opt-in.

Strengthened (User mandate, 2026-05-29): the parallel-stream floor is raised from ≥ 2 to ≥ 3 and an **auto-backfill rule** is added — the moment any one stream is FULLY done, a new stream MUST immediately start and take its place, so the count NEVER drops below 3 while actionable workable items remain. Verbatim mandate: *“make sure there are always at all times at least 3 parallel working streams parallel with main work stream as long there are workable items we can do! as soon as one working stream is fully done, another new MUST start work and take its place.”*

(A) Main stream stays FREE — ALL commit AND push run detached, never blocking. Every commit operation AND every push operation MUST run in the background (detached via `nohup ... > <log> 2>&1 & + disown`, per §11.4.88’s commit-lock-release-immediately + detached-push seam). The conductor’s main work stream MUST NOT block on a commit landing or a push completing — the moment the local commit is durable, the main stream returns to the priority queue while the push fans out to every upstream in the background. Blocking the main stream on a push (or a slow mirror) is a §11.4.103 violation (the exact 2026-05-26 §11.4.88 forensic harm — a 5-hour flock hold — generalised to the routine layer).

(B) ≥ 3 parallel background streams at all times whenever actionable work exists + auto-backfill. Whenever the queue holds three or more non-contending actionable items, the conductor MUST run **at least three** parallel background streams alongside the main stream (subagent-driven per §11.4.70 / §11.4.20, isolated per the §11.4.58 PWU worktree model + §11.4.84 working-tree quiescence). **Auto-backfill rule (User mandate, 2026-05-29):** the moment any one stream is FULLY done (its workable item closed with anti-bluff evidence per sub-rule (E), or merged), a new stream MUST IMMEDIATELY start and take its place by claiming the next-highest-priority non-contending item from the queue — the active-stream count NEVER drops below 3 while actionable workable items remain. Dropping below 3 active streams while three-plus disjoint actionable items wait — or serialising work that could run in parallel — is a §11.4.103 violation. The count may fall below 3 ONLY when there genuinely are no remaining non-contending actionable workable items OR every remaining item is externally blocked (per §11.4.94 / §11.4.97 / §11.4.101 idle-only-when-blocked). The §12.6 60% memory ceiling and the §11.4.58 6-concurrent-agent cap bound the stream count from above (so the standing band is 3–6 background streams).

(C) Most-critical + most-visible first; audio always top. Streams are filled in strict priority order — most-critical and most-user-visible items first — and **audio fixes are always the apex priority** per §11.4.72. When the conductor faces a choice of which item to dispatch onto a free stream, the highest-severity / most-user-visible / audio item wins; research, refactors, and low-stakes tasks fill remaining capacity only after the critical set is dispatched (per §11.4.42 priority-ordered iteration).

(D) Safe-during-build scope only while a heavy build runs. While the 42 GB containerised AOSP rebuild (§12.9) — or any heavy `gradle / m -j` build — is in flight, the parallel background streams **MUST** restrict themselves to the §11.4.96 SAFE catalogue: investigation / forensic analysis, documentation, test-authoring, gate + paired-mutation authoring, read-only on-device probes, submodule edits, web research, workable-items DB ops, and backgrounded pre-build / meta-test runs. The streams **MUST NOT** launch a second `gradle / m -j` / containerised build concurrent with the running one (§12.8 concurrency-hardening — a second heavy JVM/Soong build storms the host and risks the §12.7 OOM cascade). **UNSAFE-during-build operations** (§11.4.96 (α)-(θ)) are deferred to between-build windows.

(E) Heavy anti-bluff on every closure. Every item closed by a parallel stream **MUST** carry real captured evidence and survive the anti-bluff covenant family — root cause proven or explicitly marked **UNCONFIRMED: / UNKNOWN: / PENDING_FORENSICS:** per §11.4.6, captured-evidence per §11.4.5 + §11.4.69, deterministic consistency per §11.4.50, paired §1.1 meta-test mutations, and the structured root-cause investigation of §11.4.102 the moment any problem surfaces. A stream that closes an item on metadata-only / configuration-only / absence-of-error / grep-without-runtime evidence is a §11.4 PASS-bluff regardless of how parallel the routine looked.

(F) Idle ONLY when genuinely externally blocked. The conductor idles **ONLY** when every queued item is genuinely blocked on an external result it cannot accelerate (hardware access, network upstream, an in-flight build / test / push completion) **OR** the operator issues **STOP OR §12 host-session-safety demands it** (per §11.4.94(A) + §11.4.97). “Nothing visible to do” with progressable non-contending items in the queue is **NEVER** a valid idle reason. Posing a question and going idle while non-blocked items remain is a §11.4.94 + §11.4.97 + §11.4.103 violation; an unavoidable block parks one work unit, not the whole loop (§11.4.101 maximize-progress-while-blocked).

Composes with §11.4.58 (parallel-development PWU pipeline — §11.4.103 is its standing-routine framing), §11.4.70 / §11.4.20 (subagent-driven default — the parallel streams **ARE** subagents), §11.4.72 (audio top-priority — sub-rule (C)), §11.4.87 (endless-loop autonomous work — §11.4.103 is the always-on operating mode that anchor governs), §11.4.88 (background-push — sub-rule (A) is its routine-layer generalisation), §11.4.89 (background test execution — backgrounded tests are one stream class), §11.4.94 (zero-idle priority-first parallel-by-default — sub-rules (B)+(C)+(F) bind it), §11.4.96 (safe-parallel-work-with-long-build catalogue — sub-rule (D) cites it), §11.4.97 (maximum-use-of-idle-time — sub-rule (F)), §11.4.101 (autonomous-decision-over-blocking — maximize-progress-while-blocked), §11.4.102 (systematic-debugging on any spotted problem inside any stream), §11.4.42 (priority-ordered iteration), §11.4.84 (working-tree quiescence for concurrent subagent commits), §12.6 / §12.7 / §12.8 / §12.9 (host-session safety bounds stream count + forbids a second heavy build), §9.2 (data-safety for any destructive op a stream would attempt).

Classification: universal (§11.4.17) — continuous parallel-stream working with a free main stream, priority-first dispatch, and anti-bluff closures is a reusable operating discipline for ANY project on ANY runtime that supports background execution + subagent delegation; nothing in the routine is hardware-, vendor-, or region-specific.

4-layer coverage per §11.4.4(b). Propagation gate CM-COVENANT-114-103-PROPAGATION enforces the literal anchor 11.4.103 across the canonical consumer fleet (parent + owned-submodule CLAUDE.md / AGENTS.md / QWEN.md). Paired §1.1 meta-test mutation strips the 11.4.103 literal from a consumer file → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md §11.4.103 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no --block-main-stream, --synchronous-commit, --synchronous-push, --single-stream-only, --skip-parallel-streams, --serialise-actionable-work, --idle-without-queue-survey flag exists. The 2026-05-29 user mandate is unambiguous: continuous parallel-stream working with a free main stream is the regular working routine, not a per-request opt-in.

§11.4.104 — Participant identity, attribution & notification-tagging (User mandate, 2026-05-31)

Short tag: participant-attribution-tagging.

Forensic anchor — verbatim user mandate (2026-05-31):

“Every supported messenger must relate messages to participants (Subscribers/Users); the same logical person may have a different username on every messenger. Workable items must carry who created them and who they are assigned to. Notifications must @-tag the right participant — but never the operator (who drives the system) and never the system agent.”

This anchor binds participant identity, workable-item attribution, and notification @-tagging as MANDATORY constraints on every consumer that ships a messenger/notification surface (Herald and its flavor binaries are the reference implementation; ATMOSphere and any future messenger-bearing project inherit per §11.4.35). The detailed normative spec is the Herald design contract docs/design/PARTICIPANT_ATTRIBUTION.md (model, columns, attribution rules, tagging matrix) — this section restates its load-bearing constraints; implementations code against the contract, not against a paraphrase.

(A) Participant identity — logical person + per-channel handle. Every supported messenger MUST relate inbound + outbound messages to a **Participant** (a logical Subscriber/User). A single logical person/agent MAY carry a DIFFERENT username on every messenger; identity is therefore modelled as a logical subscriber (canonical, messenger-neutral handle, kind ∈ {human, agent, service}) PLUS per-channel aliases (channel, channel_user_id, the per-channel @username used for tagging). The **canonical handle** is the string stored in

attribution columns — the closed set is Claude (the reserved system-agent sentinel; kind=agent) OR a human's canonical handle (a subscriber @username, messenger-neutral, resolved per-channel via the alias table).

(B) Operator designation — env var, not a DB flag. The **Operator** is the one human who drives the system via the agent CLI. The Operator is designated by the environment variable `HERALD_<CHANNEL>_OPERATOR_USERNAME` (e.g. `HERALD_TGRAM_OPERATOR_USERNAME`, `HERALD_SLACK_OPERATOR_USERNAME`) — per messenger, NOT a database flag. The Operator is a normal Participant whose canonical handle equals that env value.

(C) Workable-item attribution — created_by + assigned_to. Every workable item MUST carry two canonical-handle columns: `created_by` and `assigned_to`. Attribution rules: opened via the agent CLI prompt (operator-driven) → `created_by = Operator`; opened by the system/agent detecting an issue/task/improvement/missing-feature → `created_by = "Claude"`; received THROUGH the system as a subscriber message → `created_by = the sender's resolved canonical @username`. `assigned_to` **defaults to the Operator** and MAY be overridden explicitly (a prompt/message assigning to @someoneelse). Both columns store the canonical handle string and are self-contained in the SSoT + its Markdown export. Legacy items predating this anchor carry empty ("") attribution and MUST still parse + validate.

(D) Notification @-tagging matrix. On any workable-item event, the outbound notification dispatched to each messenger group @-tags the participant(s) who must be aware, resolved to that channel's @username:

```
mentions = {}
if assigned_to is a human handle AND assigned_to !=
Operator:           mentions += assigned_to
if created_by is a human handle AND created_by != Operator AND
created_by != "Claude": mentions += created_by
# "Claude" is NEVER tagged (it is the system agent).
# the Operator is NEVER tagged (drives the system; no self-ping).
# de-dup; for each mention resolve the @username on the target
channel – skip if the participant has no alias there.
```

This satisfies the operator's stated cases exactly: assigned-to-Operator → no tag; opened-by-Operator-assigned-to-another → tag the assignee; opened-by-a-non-Operator-non-Claude-subscriber → tag the creator.

(E) Anti-bluff (MANDATORY, composes with §11.4 / the end-user quality covenant). Every layer ships unit + integration + E2E + full-automation tests producing REAL captured evidence (no metadata-only / absence-of-error PASS, no false positive/negative): real SQLite round-trip with the new columns (byte-identical, including legacy fixtures that DON'T carry the fields); the tagging matrix proven by a truth-table test with a mutation that flips one cell forcing a FAIL; an E2E real-item-event → real-dispatched-message asserting the exact @usernames AND a NEGATIVE case proving the Operator is NOT tagged. Evidence committed under docs/qa/<run-id>/.

Composes with §11.4 + §11.4.1..§11.4.16 (end-user quality / anti-bluff covenant — sub-rule (E) is bound by it), §11.4.5 / §11.4.69 (captured evidence), §11.4.50 (deterministic consistency), §11.4.91 (workable-item field clarity), §11.4.93 / §11.4.95 (the SQLite workable-items SSoT the created_by/assigned_to columns extend), §1.1 (paired-mutation proof of the tagging matrix).

Classification: universal (§11.4.17) — participant identity, item attribution, and notification @-tagging is a reusable discipline for ANY project that ships a messenger/notification surface with a workable-items registry; the env-var operator designation + canonical-handle model + tagging matrix are vendor-neutral. Projects with NO messenger surface inherit the anchor latently (it binds the moment they ship one) — the §11.4.96 “principle binds even absent the surface” restatement pattern.

4-layer coverage per §11.4.4(b). Propagation gate CM-COVENANT-114-104-PROPAGATION enforces the literal anchor 11.4.104 across the canonical consumer fleet (parent + owned-submodule CLAUDE.md / AGENTS.md / QWEN.md). Paired §1.1 meta-test mutation strips the 11.4.104 literal from a consumer file → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md §11.4.104 in the HelixConstitution submodule; detailed spec Herald docs/design/PARTICIPANT_ATTRIBUTION.md. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker. No escape hatch — no --skip-attribution, --no-participant-tagging, --tag-operator-anyway, --attribution-later flag exists.

§11.4.105 — Natural-language intent recognition & clarification (User mandate, 2026-05-31)

Short tag: intent-recognition-clarification.

Forensic anchor — verbatim user mandate (2026-05-31):

“Users must NOT need to know command syntax (no COMMAND: ... prefix). They send a clear natural-language message; the System determines the intent. The System recognizes the commands it has; if none matches it infers the exact intent; if it is totally unable it replies, tags the user (@user ...), and asks to clarify precisely. We MUST always do our best to determine exact intent so we never annoy end users. This is a CORE part of the System.”

This anchor binds natural-language intent recognition as a MANDATORY constraint on every consumer that ships a messenger/notification-driven command surface (Herald and its flavor binaries are the reference implementation; ATMOSphere and any future messenger-bearing project inherit per §11.4.35). The detailed normative spec is the Herald design contract docs/design/INTENT_RECOGNITION.md (three-tier resolution, command set, the clarify action, envelope wiring) — this section restates its load-bearing constraints; implementations code against the contract, not against a paraphrase.

(A) No required command syntax. Users MUST NOT be required to know any command syntax — no COMMAND: prefix, no rigid grammar. They send a clear natural-language message; the System is responsible for determining the intent. Any design that forces an end user to memorize syntax to be understood is a violation of this anchor.

(B) Three-tier intent resolution (the mandatory discipline). Every inbound subscriber message MUST be resolved to exactly one action via three tiers, in order — the first that succeeds wins:

TIER 1	command recognition – recognize the System's existing command set from	natural language and map it to a structured action	(a confident, deterministic match → that action).
TIER 2	intent inference – when NO command matches, infer the exact intent from	natural language to	the message (the LLM dispatch maps the right action) – NEVER guessing.
TIER 3	clarify (fallback) – when neither a command nor a confident intent can be	message, TAG the	determined, REPLY to the original sender (`@username`, resolved via the §11.4.104 IdentityResolver) and ask a PRECISE clarifying question that NAMES the candidate intents. No guessing, no silent drop.

Tier 3 is the anti-annoyance guarantee: the user is never ignored and never has to learn syntax — at worst they receive a friendly, specific “@user, did you mean X or Y?”.

(C) Never guess, never drop. The System MUST NEVER guess an action — a wrong action is worse than a clarifying question (composes with §11.4.6 no-guessing). The System MUST NEVER silently drop or ignore an inbound message. We always do our best to determine the exact intent so end users are never annoyed; only genuine ambiguity reaches Tier 3, and Tier 3 always replies-tags-and-asks rather than failing silently.

(D) Anti-bluff (MANDATORY, composes with §11.4 / the end-user quality covenant). Every tier ships unit + integration + E2E + full-automation tests producing REAL captured evidence (no metadata-only / absence-of-error PASS, no false positive/negative): Tier 1 proven by a truth-table of natural-language messages → expected action+fields PLUS the conservative negatives that MUST fall through to “no match”; Tier 3 proven by an E2E ambiguous-message dispatch whose recorded reply body is EXACTLY @<sender> <specific question> (the user is tagged + asked, not ignored) AND a NEGATIVE proving a clear command does NOT trigger clarify; a paired §1.1 mutation that breaks the recognizer’s confidence guard (forcing a false-match) OR drops the clarify tag MUST FAIL a test. Evidence committed under docs/qa/<run-id>/.

Composes with §11.4 + §11.4.1..§11.4.16 (end-user quality / anti-bluff covenant — sub-rule (D) is bound by it), §11.4.6 (no-guessing — a wrong action is worse than a clarifying question), §11.4.104 (participant identity — the clarify reply tags the sender via the IdentityResolver), §11.4.5 / §11.4.69 (captured evidence), §11.4.98 (full-automation — every tier’s test is self-driving), §1.1 (paired-mutation proof of the confidence guard + clarify tag).

Classification: universal (§11.4.17) — natural-language intent recognition with a three-tier recognize→infer→clarify discipline is a reusable interaction contract for ANY project that ships a messenger/command surface; the no-syntax mandate, never-guess rule, and reply-tag-and-ask fallback are vendor-neutral. Projects with NO messenger surface inherit the anchor latently (it binds the moment they ship one) — the §11.4.96 “principle binds even absent the surface” restatement pattern.

4-layer coverage per §11.4.4(b). Propagation gate CM-COVENANT-114-105-PROPAGATION enforces the literal anchor 11.4.105 across the canonical consumer fleet (parent + owned-submodule CLAUDE.md / AGENTS.md / QWEN.md). Paired §1.1 meta-test mutation strips the 11.4.105 literal from a consumer file → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md §11.4.105 in the HelixConstitution submodule; detailed spec Herald docs/design/INTENT_RECOGNITION.md. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker. No escape hatch — no --require-command-syntax, --guess-intent-ok, --skip-clarify, --drop-on-ambiguous flag exists.

§11.4.106 — Docs Chain — mechanical documentation/DB sync engine (Operator mandate, 2026-05-31)

Short tag: docs-chain-mechanical-sync.

Forensic anchor — operator mandate (2026-05-31):

Docs Chain is the canonical mechanical enforcer of the documentation-sync mandates. Consuming projects **MUST** use the engine instead of ad-hoc per-project doc-sync scripts; they register their own chains via per-context YAML; the engine never fakes a transform and every sync/verify carries real captured evidence.

Docs Chain — the vasic-digital/docs_chain engine (a universal, Go-implemented, **bidirectional document-and-database dependency-propagation engine**) — is the canonical mechanical enforcer of the family of documentation-sync mandates this Constitution already declares. This is a **MANDATORY** constraint binding every consuming project. The detailed normative spec is the engine’s own documentation set — docs/CONSTITUTION_INTEGRATION.md (distribution + inheritance + the anchor mapping table) and docs/USE_CASE_CATALOGUE.md (the ready-to-use chain recipes); this section restates its load-bearing constraints, it does not redefine them. Implementations code against the engine’s contract, not against a paraphrase.

(A) Use the engine, never ad-hoc scripts. Consuming projects **MUST** drive their documentation/DB sync through Docs Chain — consumed **by reference** via the flat-layout sibling (the `~/Projects/docs_chain` working tree during local development / the path the constitution exposes once the Phase-6 submodule pointer lands), inherited the same way §11.4.80 inherits the `codegraph_*` scripts: referenced at the constitution-exposed path, **NEVER copied** into the consumer. Ad-hoc per-project doc-sync scripts (`sync_issues_docs.sh`, `generate*_summary.sh`, `sync*_export.sh`, `update_readme_doc_links.sh`, et al.) are superseded by the engine and **MUST** be retired as each matching context is registered.

(B) Consumer-owned contexts via config. The engine is project-agnostic; it carries NO project-specific path or chain. The consumer registers its chains as data via `.docs_chain/contexts/*.yaml` at its own project root — one independent context per file — per the §11.4.28 decoupling rule (the submodule never carries project-specific context; the consumer injects its chains via config). `state.json` (at `<project-root>/docs_chain/state.json`, regenerated on demand by `docs_chain sync`) and the `*.docs_chain.tmp` staging sidecars are **gitignored**; the consumer owns those `.gitignore` entries.

(C) The sync anchors it mechanizes. Docs Chain is the mechanical engine behind a family of existing documentation-sync anchors; once a consumer registers the matching context it enforces the anchor **in place of the ad-hoc script** — see the `docs/CONSTITUTION_INTEGRATION.md` mapping table: §11.4.12 (Issues_Summary always-sync), §11.4.53 (Fixed_Summary parity), §11.4.45 (Status.md maintenance), §11.4.56 (Status_Summary two-audience parity), §11.4.57 (README doc-link section), §11.4.59 (README always-sync export), §11.4.60 (documentation composite-sync), §11.4.65 (universal markdown export), §11.4.86 (roster/corpus auto-sync), §11.4.93 (workable-items DB single source of truth), §11.4.95 (DB tracked + WAL-checkpoint commit), §12.10 (CONTINUATION maintenance), and §11.4.44 (document revision header — exports kept in sync). Engine-wide guarantees: change detection by **content hash, NOT mtime** (§11.4.86); **atomic-rename + SQLite-transaction commit** with rollback on any transform error and optional hardlink backup (§9.2 zero-risk data safety); a both-dirty sync pair surfaces a **conflict** — **never a silent merge** (exit 2, §11.4.6 no-guessing); `verify` is the deterministic CI / pre-build sink-side gate over byte-stable transforms (§11.4.50); per-run captured evidence lands at `qa-results/docs_chain/<run-id>/` (§11.4.69).

(D) It does NOT replace authoring discipline. Docs Chain replaces only the **mechanical sync** an anchor requires — it does NOT replace the authoring discipline of any anchor. The source author still writes the §11.4.44 revision header into the source; the engine then keeps every derived export in sync. A consumer must not treat the engine as licence to skip authoring.

(E) Anti-bluff (MANDATORY, composes with §11.4 / the end-user quality covenant). Docs Chain **NEVER** fakes a transform: a missing `pandoc` / `weasyprint` (or any absent tool) surfaces a typed `ToolAbsentError (IsToolAbsent)` and an honest §11.4.3 SKIP-with-reason — **never a fake PASS**, never a partial write. Every sync / `verify` run carries real captured evidence (the `qa-results/docs_chain/<run-id>/` artefact). A metadata-only / “absence-of-error” / config-only PASS at the sync layer is a §11.4 PASS-bluff.

Status note (§11.4.6 no-guessing). The engine's Phases 1–3 (core DAG + content-hash + adapters + atomicity/orchestrator) AND the cmd/ CLI / per-context YAML loader (Phase 4) are IMPLEMENTED + tested — cmd/docs_chain/main.go registers the sync / verify / diff (alias of verify) / doctor subcommands wiring the config loader → state → orchestrator, and go test -count=1 ./... is GREEN at 7/7 packages (incl. internal/orchestrator spec_wiring_test, both PASS). FACT (§11.4.6): this CLI/loader lives in the working tree but is UNTRACKED (git ls-files docs_chain/cmd/docs_chain/main.go is empty) and the engine is NOT yet a registered git submodule (git submodule status shows no docs_chain entry; in-tree at ./docs_chain, HEAD = parent HEAD). The constitution-submodule distribution (Phase 6) remains PLANNED + OPERATOR-GATED. This anchor states the binding contract; consumers wire the engine as each phase lands, and MUST NOT claim working behaviour a phase has not yet shipped.

Classification: universal (§11.4.17) — a content-hash bidirectional doc/DB propagation engine is a vendor-neutral mechanism reusable by ANY governed project. Projects with no derived-export or DB-sync surface inherit the anchor latently (it binds the moment they ship one) — the §11.4.96 “the principle binds even absent the surface” restatement pattern.

4-layer coverage per §11.4.4(b). Propagation gate CM-COVENANT-114-106-PROPAGATION enforces the literal anchor 11.4.106 across the canonical consumer fleet (parent + owned-submodule CLAUDE.md / AGENTS.md / QWEN.md). Paired §1.1 meta-test mutation strips the 11.4.106 literal from a consumer file → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Composes with §11.4 + §11.4.1..§11.4.16 (end-user quality / anti-bluff — sub-rule (E) is bound by it), §11.4.6 (no-guessing — both-dirty sync → conflict, never silent merge), §11.4.12 / §11.4.45 / §11.4.53 / §11.4.56 / §11.4.57 / §11.4.59 / §11.4.60 / §11.4.65 / §11.4.86 / §11.4.93 / §11.4.95 / §12.10 / §11.4.44 (the documentation-sync anchors it mechanizes), §11.4.28 (engine/context decoupling), §11.4.80 (inherited-by-reference, never-copied — the codegraph pattern), §9.2 (atomic commit + rollback), §11.4.50 (deterministic verify gate), §11.4.69 / §11.4.5 (captured evidence), §1.1 (paired-mutation propagation proof).

Canonical authority: this Constitution.md §11.4.106 in the HelixConstitution submodule; detailed spec the Docs Chain engine docs docs/ CONSTITUTION_INTEGRATION.md + docs/USE_CASE_CATALOGUE.md (vasic-digital/docs_chain). All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker. No escape hatch — no --ad-hoc-sync-ok, --skip-docs-chain, --fake-transform, --sync-evidence-optional flag exists.

§11.4.107 — Anti-bluff AV/test-validation techniques mandate (User-driven research, 2026-06-02)

Short tag: av-liveness-oracles.

Forensic anchor — captured incidents (2026-06-02, genericised): a test asserting “media is playing on the target output” PASSED on the strength of a SINGLE captured frame that showed “a picture” — but the picture was a FROZEN / STALE frame left over from the previously-played content (a stale-producer / stuck-decoder failure mode), so the feature was broken for the end user while the test was green. A sibling incident: media briefly FLASHED on the WRONG output (the source/primary surface) for ~1 s before routing to the intended output, and no test sampled the wrong output during the transition window, so the glitch shipped undetected. A third class: a comparator (color / freeze / OCR oracle) that was itself broken — it PASSED on a deliberately-degraded artefact — meaning the analyzer, not just the feature, was the bluff. §11.4.5 mandates *captured* evidence and a baseline presence pass; §11.4.107 raises the bar to **liveness + correct-routing + self-validated-analyzer** so a single still frame, a wrong-output flash, or a broken oracle can no longer pass.

The mandate. Every test asserting that audio or video output is genuinely playing / advancing for the end user MUST satisfy ALL of the following — each is a §11.4 PASS-bluff to omit:

1. **A single captured frame is NOT proof of playback.** Prove LIVE, ADVANCING frames over a steady-state window ($\geq N$ frames across $\geq$ a few seconds) via a **freeze-detection oracle** — a near-duplicate detector (the platform’s frame-difference / freeze-detect-class filter OR a perceptual-hash adjacent-frame distance), **NOT a byte-identical compare** (a re-encoded / 1-pixel-noisy stale frame defeats byte-identity; keep byte-identity only as a zero-cost early-out pre-filter). Adjacent frames MUST differ above a calibrated threshold; $\geq$ a sustained interval of near-identical frames $\Rightarrow$ FROZEN $\Rightarrow$ FAIL.
2. **An independent frame-advance counter from the platform’s compositor / decoder telemetry** MUST increase across the window — an oracle in a DIFFERENT domain from the pixel check (compositor frame-timestamp advance, decoder `framesDecoded` / `framesRendered`, or equivalent). A flat counter over the window $\Rightarrow$ stuck decoder $\Rightarrow$ FAIL, even if the pixels appear to move. Keeping a pixel-domain oracle AND a telemetry-domain oracle is metamorphic redundancy: a coordinated false-PASS now requires two independent failures.
3. **Loading / buffering is a DISTINCT state — wait for genuine playback-start before judging liveness.** Never false-FAIL a still-loading stream, never false-PASS a loading spinner. Poll the decoder frame-advance counter / first-advancing-frame / buffering flag until playback genuinely starts (generous per-class timeout). Timeout + content endpoint unreachable $\Rightarrow$ SKIP-with-reason per §11.4.3 (`network_unreachable_external` / `geo_restricted`); timeout + reachable $\Rightarrow$ FAIL with captured evidence. Only AFTER confirmed playback-start run the liveness checks over the steady-state window.
4. **Not-stale-from-previous cross-check.** Before asserting new content N is live, confirm its first frame DIFFERS from the previous content’s last captured frame (perceptual-hash mismatch). Identical $\Rightarrow$ the “new” content is a leftover frozen frame $\Rightarrow$ FAIL. Critical for switch-between-items sequences.
5. **Measured FPS / no-lost-frames within tolerance** of the source’s expected rate (where the source rate is knowable) or a sane floor otherwise; drop ratio above budget $\Rightarrow$ FAIL (§11.4.5 frame-health).

6. **No-flash-on-the-wrong-output during a routing transition.** When content is meant to appear ONLY on a target output, sample the NON-target output at high frequency (e.g. every ~200 ms) across the transition window; ANY sampled wrong-output frame containing the content $\Rightarrow$ FAIL (brief-flash glitch). (Where a content-protection regime legitimately forces a particular output, that regime is the expected case, not a fail — classify it explicitly, never by guessing per §11.4.6.)
7. **Drive tests through the realistic user path (feed / UI), not deep-link shortcuts.** Routing / transition / alternation bugs manifest on the path a real user takes (selecting from a feed / list via the UI); deep-link / direct-dispatch shortcuts bypass the very transition paths where these bugs live. Where the UI is not introspectable, fall back AND flag operator-attended per §11.4.52 — never fake-PASS. (This is the §11.4.48 UI-driven discipline applied to the liveness battery.)
8. **The oracle problem — use metamorphic relations when there is NO golden source.** For content the project does not own (streamed media), there is no golden reference, so full-reference comparison is impossible. Encode metamorphic relations instead: the SAME content captured on output-A vs output-B must be structurally equivalent (after resolution normalisation); a paused stream's frame-advance counter must STOP; a 2 $\times$ speed must roughly double the frame-advance rate; the same clip played twice must produce equivalent liveness metrics. Each is a relation, not a fixed expected value — robust to content variety.
9. **Full-reference quality metrics vs a golden source for content the project DOES own.** Where a source reference exists (local-file playback of owned media, a test corpus), assert perceptual-quality / color fidelity (SSIM / VMAF / ΔE_{2000} -class) against the temporally-aligned source so a pale / desaturated / wrong-range / soft-scaled render is caught across motion, not just at one sampled frame.
10. **Mutation-test every analyzer with a golden-good + golden-bad fixture pair** so the analyzer itself provably cannot bluff. For every comparator (freeze / motion / color / sharpness / aspect-ratio / FPS / OCR / audio-presence) keep a clean control fixture (must PASS) AND a deliberately-degraded fixture (must FAIL — frozen clip, black-gap clip, desaturated clip, downsampled clip, stretched clip, half-rate clip, overlay-baked frame, silent track). An analyzer that PASSES its golden-bad fixture is a bluff gate. This is the §1.1 paired-mutation discipline applied to the AV *analyzers*, not only the grep gates. Document any behaviourally-equivalent mutant expected NOT to flip so a survivor there is not a false alarm.
11. **Per-channel audio integrity, not a single aggregate number.** A single aggregate RMS misses a dead channel in a multichannel stream and a near-silent / DC-offset stream that scrapes past a bare floor. Assert per-channel RMS / loudness (an EBU-R128-class integrated/momentary loudness + per-channel level + silence-region census + DC-offset / NaN-Inf check); EVERY declared channel above floor. On-device, take an underrun / XRUN census from the audio subsystem — a starved pipeline is a defect even if the captured waveform “has sound”.

12. **OCR overlay / subtitle detection needs a per-word confidence floor + a region-of-interest** to avoid BOTH a false-positive (noise read as a hostile overlay $\Rightarrow$ false-FAIL of a working feature) AND a false-negative (a real overlay missed $\Rightarrow$ false-PASS). Discard sub-confidence words; OCR the ROI where the text lives (bottom-third for subtitles, centre for dialogs); emit per-finding confidence so a reviewer sees WHY a finding fired. Where a subtitle must appear on one output and NOT another, assert both sides (a property relation per point 8).
13. **Thresholds MUST be calibrated on the project's own fixtures, not hardcoded from literature (§11.4.6 no-guessing).** Every freeze / motion / SSIM / VMAF / loudness / OCR-confidence threshold is recorded as captured evidence derived from a known-good and known-bad sample on the project's own capture path — never a literature constant asserted as fact. The capture path's own re-encode loss is part of the baseline.

Honest gaps (§11.4.6). True physically-displayed (photon) frame rate has no clean software oracle — the compositor / decoder counters measure the *presentation pipeline*, not photons at the sink; a high-speed external camera rig is the only ground truth. The counters are sufficient to catch a stuck / half-rate decoder and MUST be used, with the photon-FPS gap flagged honestly rather than claimed. Broadcast-QC vendor tooling validates that freeze / duplicate / lost-frame / black detection is industry-standard practice, but the exact proprietary thresholds are undocumented; the project adopts the published open-method analogues and calibrates per point 13.

4-layer coverage per §11.4.4(b). Per-fix: pre-build gate (e.g. a CM-AV-LIVENESS-NO-FROZEN-FRAME-class gate asserting every test that claims output-is-playing references the liveness battery — freeze oracle + frame-advance counter + not-stale cross-check + FPS — not a single-frame check; plus an analyzer-self-validation gate wiring the golden-good/golden-bad fixtures into the project's meta-test) + on-device / runtime test + **paired §1.1 meta-test mutation** (a test asserting output-is-playing with ONLY a single-frame check $\rightarrow$ gate FAILs; an analyzer that PASSes its golden-bad fixture $\rightarrow$ self-validation FAILs) + HelixQA Challenge entry for the user-visible feature. Every such PASS is emitted via `ab_pass_with_evidence` citing the motion / not-stale / frame-advance / fps / per-channel-loudness / metamorphic / self-validation artefacts (feature classes `video_display`, `audio_output`, `subtitle_render` per §11.4.69) — NEVER a single screenshot.

Classification: universal (§11.4.17) — freeze-detection oracles, frame-advance telemetry, loading-aware playback-start waits, not-stale cross-checks, metamorphic relations, full-reference quality metrics, golden-good/golden-bad analyzer self-validation, per-channel loudness, and confidence-floored OCR are platform-neutral AV / test-validation techniques reusable by ANY project that validates media playback or any pixel/audio output; the project supplies the concrete capture mechanism (its compositor / decoder telemetry, its sink-introspection API per §11.4.13, its capture tool) and the calibrated thresholds per §11.4.35.

Composes with §11.4.5 (strict expansion of its video + audio quality bar to liveness + correct-routing), §11.4.6 (no-guessing — thresholds calibrated + captured, loading-vs-frozen disambiguated by evidence, content-protection regime classified not guessed), §11.4.50 (deterministic-consistency — the liveness battery runs N iterations identically), §11.4.68 (positive sink-side / downstream evidence — the frame-advance counter / sink-introspection cross-check is its video analogue), §11.4.69 (universal sink-side positive-evidence taxonomy — `video_display` / `audio_output` / `subtitle_render` classes), §11.4.85 (stress + chaos — liveness

under sustained load + transition chaos), plus §11.4.2 (recorded evidence), §11.4.3 (topology SKIP-with-reason for unreachable content), §11.4.13 (sink-side introspection as one oracle), §11.4.48 (UI-driven path), §11.4.52 (autonomous validation; operator-attended fallback only when UI not introspectable), §1.1 (paired-mutation discipline applied to the analyzers).

Propagation. Propagation gate CM-COVENANT-114-107-PROPAGATION enforces the literal anchor 11.4.107 across the consumer fleet (parent + owned-submodule CLAUDE.md / AGENTS.md / QWEN.md); paired §1.1 meta-test mutation strips the 11.4.107 literal from a consumer file → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md §11.4.107 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker. No escape hatch — no --single-frame-proves-playback, --skip-liveness, --byte-identical-freeze-OK, --no-frame-advance-counter, --skip-not-stale-check, --allow-wrong-output-flash, --deep-link-shortcut-OK, --unvalidated-analyzer-OK, --aggregate-rms-suffices, --hardcoded-thresholds-OK flag exists.

§11.4.108 — Four-layer fix-verification + runtime-signature-as-definition-of-done mandate (systematic-debugging Phase 4.5, 2026-06-03)

Short tag: runtime-signature-definition-of-done.

Forensic anchor (genericised, 2026-06-03). Across one batch, multiple fixes were “green” at every gate yet NOT working for the end user: a change present in the source tree and passed by both the pre-build (source) gate and the post-build gate never reached the boot-time kernel command-line embedded in the deployed boot artifact, so the corresponding output feature was dead despite all-green; sibling fixes correctly built into the system image were masked by stale per-user overlay copies left from a previous deployment (the running code was the stale overlay, not the freshly-built system code), so they merely *looked* broken; deployment did not produce a clean state (the mutable overlay was not wiped), so the stale shadow survived; and the validation suite ran against whatever was running — which could be the stale shadow — so it reported green on code that was never actually exercised. Per superpowers:systematic-debugging Phase 4.5: when each fix reveals a fresh “fixed-but-not-working” in a *different* place, that is NOT a series of independent bugs — it is ONE architectural VERIFICATION flaw, and patching each symptom is thrashing.

The architectural flaw (FACT, traced). A fix must cross FOUR distinct layers, and “fixed” at one layer does NOT imply fixed at the next: (1) **SOURCE** — the change is committed in the source file (what a grep-the-source pre-build gate checks); (2) **ARTIFACT** — the change’s BYTES actually landed in the produced build artifact (image / bundle / installer / boot artifact / embedded command-line); (3) **RUNTIME-ON-CLEAN-TARGET** — the change is active on a *clean / fresh* deployment, the layer the end user experiences, with no stale overlay shadowing the deployed code; (4) **USER-VISIBLE** — the feature actually works for the end user (the §11.4.5 / §11.4.69 captured-evidence layer). Green at layer 1 is the cheapest,

least conclusive signal; nothing being green at layer 1 (or even layer 2) tells you about layers 3–4. A gate that verifies ONLY the source layer is insufficient and is itself a §11.4 bluff surface.

The mandate (ALL must hold).

1. **Runtime-signature-as-definition-of-done.** A fix is DONE only when its declared **runtime signature** is verified on a CLEAN / fresh deployment (the layer the end user experiences). Source-committed $\neq$ artifact-contains-it $\neq$ active-on-clean-target $\neq$ user-visible-working — these are FOUR distinct claims; proving one does NOT prove the next.
2. **Every fix declares ONE machine-checkable runtime signature** — a single observable on a clean target that proves the fix is BOTH active AND working (a property read from the running system, a downstream/sink-side report per §11.4.13, a captured-evidence assertion per §11.4.5 / §11.4.69, or a counter/state delta — never a re-grep of the source). The **registry of these per-fix runtime signatures is the SINGLE SOURCE OF TRUTH for “fixed”** — it REPLACES “a gate greps the source” as the definition of done.
3. **Verification gates MUST span all four layers**, per fix: **source** (pre-build / pre-merge), **artifact** (post-build — assert the change’s BYTES landed in the produced artifact, NOT merely that the source still contains the change; cf. the embedded-command-line / stripped-asset failure mode), **runtime-on-clean-target** (post-deploy — assert the runtime signature on a freshly-deployed clean target), **user-visible** (captured evidence per §11.4.5 / §11.4.69). A gate that asserts only the source layer is insufficient and is itself a §11.4 bluff surface.
4. **Eliminate the stale-deployment / shadow layer — by construction.** Deployment MUST yield a CLEAN state (wipe the mutable overlay / per-user mutable state that can shadow the deployed build) **OR** a pre-validation assertion MUST prove running-artifact == built-artifact (no stale override masks the freshly-deployed build) BEFORE any validation runs. Validation that runs against possibly-stale deployed state is INVALID and any PASS it produces is a §11.4 PASS-bluff — the test exercised code that was never actually deployed.
5. **Meta-rule (composes §11.4.102 Phase 4.5).** ≥ 3 “fixed-but-not-working” discoveries within a single cycle are the signal of an architectural VERIFICATION flaw, NOT three independent product bugs. On the 3rd such discovery the agent MUST STOP patching symptoms (per §11.4.4 test-interrupt-on-discovery), fix the VERIFICATION pipeline itself (close whichever of the four layers is unverified), and re-certify EVERY item in the batch through the corrected pipeline on a clean target.
6. **A batch is “validated” only after COMPREHENSIVE per-item runtime-signature verification on a clean baseline** — NOT after the batch-touched items’ own tests pass. Spot-validating only the touched items (per §11.4.40’s full-suite reasoning) misses the cross-layer gap: a touched item can be source-green while its runtime signature is dead on a clean target.

Classification: universal (§11.4.17) — the four-layer source→artifact→runtime-on-clean-target→user-visible verification gap, the per-fix runtime-signature registry as the definition of done, the artifact-byte-presence gate, the clean-deployment-or-equal-artifact pre-validation assertion, and the ≥ 3 -discoveries-means-architectural-

flaw meta-rule are platform-neutral verification-pipeline disciplines reusable by ANY project that builds an artifact, deploys it, and validates the deployed result; the consuming project supplies the concrete artifact format, its clean-deployment / equal-artifact mechanism, and its per-fix runtime-signature observables per §11.4.35.

Composes with §11.4.1 (FAIL-bluffs forbidden), §11.4.2 (recorded evidence — the runtime signature IS the recorded evidence), §11.4.4 (test-interrupt-on-discovery + retest-from-clean-baseline — clause 5’s STOP-and-fix-the-pipeline is its §11.4.108 specialisation; “clean baseline” is exactly clause 4’s clean deployment), §11.4.5 (captured-evidence quality — the user-visible layer), §11.4.6 (no-guessing — every layer asserted by captured evidence, never assumed to have propagated), §11.4.27 (no-fakes-beyond-unit — runtime signatures exercise the real deployed system), §11.4.40 (full-suite retest — §11.4.108 adds the cross-layer per-item runtime-signature dimension clause 6 demands), §11.4.46 (validate-recent-work-before-full-suite — the clean-baseline / equal-artifact pre-flight is its §11.4.108 pre-condition), §11.4.50 (deterministic consistency — the runtime signature verifies identically across N iterations), §11.4.52 (autonomous-validation — runtime signatures are machine-checkable without an operator), §11.4.69 (universal sink-side positive-evidence taxonomy — a runtime signature is a taxonomy-class observable), §11.4.102 (systematic-debugging — Phase 4.5 architectural-flaw recognition IS clause 5’s trigger).

Propagation. Propagation gate CM-COVENANT-114-108-PROPAGATION enforces the literal anchor 11.4.108 across the consumer fleet (parent + owned-submodule CLAUDE.md / AGENTS.md / QWEN.md); paired §1.1 meta-test mutation strips the 11.4.108 literal from a consumer file → the gate FAILs. A recommended per-fix gate CM-RUNTIME-SIGNATURE-REGISTRY asserts every fix declares a runtime signature verified at the artifact + runtime-on-clean-target layers (not only the source layer); paired §1.1 mutation downgrades a fix’s verification to source-only → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md §11.4.108 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker. No escape hatch — no --source-green-is-done, --skip-artifact-byte-check, --validate-against-running-state, --no-clean-deployment, --skip-runtime-signature, --spot-validate-touched-only flag exists.

§11.4.109 — Mandatory Anti-Forgetting Enforcement: PreToolUse Guard Hook + Subagent Constitutional Preamble + Orchestrator Pre-Action Checklist (Operator mandate)

Short tag: anti-forgetting-enforcement.

Forensic anchor — operator mandate:

UNCONFIRMED: pending operator’s verbatim anti-forgetting mandate quote (tracked: PENDING item 3 in the current session’s CONTINUATION.md handoff).

Background context (UNCONFIRMED quote omitted; factual description only): the motivating incident is documented in docs/AGENT_GUARDRAILS.md and in scripts/hooks/guard-forbidden-commands.sh’s in-source header. During an on-device-API build, emulator subagents ran raw host-direct emulator/adb commands instead of routing through the Containers submodule as required by the governing clauses (§6.X / §6.V / §6.AG in the consuming-project layer). The root cause was that the orchestrator **forgot to inject the rule into subagent prompts**. A rule the orchestrator forgets to paste is not enforcement — it is a social contract with the agent’s memory, and that contract is broken by every cold-session start, context-window exhaustion, and sub-agent dispatch. The mechanical fix is twofold: (1) a **PreToolUse guard hook** that blocks the forbidden command classes at the tool-call boundary unconditionally, and (2) a **canonical subagent preamble document** that the orchestrator pastes verbatim into every dispatch, plus a **pre-action checklist** for the orchestrator’s own actions. Together these are the “anti-forgetting” enforcement layer: the hook is the floor (commands cannot be executed even if the agent forgot every rule), the preamble is the ceiling (the full live ruleset, not just what the hook can pattern-match). This two-layer model was introduced in the Lava project (commit family on master, 2026-06-01/02) and is explicitly designed to be universal — scripts/hooks/guard-forbidden-commands.sh carries Classification: universal in its header and is portable by design (no project-specific paths, no jq hard dependency).

Rule. Every HelixConstitution-consuming project MUST deploy the §11.4.109 anti-forgetting enforcement layer consisting of three components, all maintained together:

(A) PreToolUse guard hook — the mechanical floor.

A shell script MUST be wired as a PreToolUse hook in the AI agent runtime’s project-scoped settings file (e.g. .claude/settings.json for Claude Code) to intercept every Bash tool call before execution and BLOCK the forbidden command classes at the tool-call boundary. The hook:

- 1. Reads the tool invocation as JSON from stdin (.tool_name + .tool_input.command).
- 2. Exits 0 for non-Bash tools (pass-through) and empty/missing commands.
- 3. For Bash tool calls, pattern-matches against the following mandatory blocked classes, each mapped to its governing clause:

Blocked class	Governing clause(s)	Examples
Raw host-direct Android emulator launch / APK install / instrumentation run	§11.4.76 (Containers mandate) + any project- layer emulator clause (e.g. §6.X / §6.V / §6.AG in Lava)	emulator -avd ..., \$ANDROID_HOME/ emulator/emulator ..., adb install ..., adb -s ... install ..., am instrument ...
Force-push / hook bypass / signature bypass	§11.4.41 (pre-force- push merge-first) + §11.4.71 (pre-push fetch) + project-layer §6.T.3 equivalent	git push --force, git push -f, git push --force-with- lease, --no-verify, --no-gpg-sign

Blocked class	Governing clause(s)	Examples
Privilege escalation	§11.4 no-sudo (project-layer §6.U equivalent)	sudo ..., su, su -, su -l ... systemctl suspend/hibernate/poweroff/reboot/halt/kill-user/kill-session, loginctl ..., pm-suspend/hibernate, shutdown ...
Host power management	§12 host-session-safety (all consuming projects)	

- Exits 2 (block) when a match is found; the stderr text printed is fed back to the agent as the refusal reason, citing the governing clause.
- Supports a documented-exception escape hatch: a command containing the literal marker `# guardrails:allow <reason>` is WARNED (stderr) but NOT blocked — EXCEPT for host-power commands, which are categorically non-overridable.

The **canonical implementation** of this hook lives at `constitution/scripts/hooks/guard-forbidden-commands.sh` in this constitution submodule. Consuming projects MUST reference it at that path (inherited by reference per §11.4.80) — NEVER copy it locally (a copy diverges silently). The script is portable: no jq hard dependency (built-in awk fallback), no consumer-project-specific paths, no hardcoded credentials.

Anti-bluff requirement: a hermetic test suite (at minimum 20 cases) MUST verify every blocked class exits 2, every allowed command exits 0, the escape hatch fires for non-power classes, and the host-power class rejects even with the escape marker. Paired §1.1 mutation: remove the `emulator -avd` pattern from the hook → the `emulator-gate` case exits 0 → test FAILs → restore → test PASSes.

(B) Subagent Constitutional Preamble — the semantic ceiling.

A canonical document MUST exist at `docs/AGENT_GUARDRAILS.md` (or the project's equivalent governance preamble path, recorded in the project's `CLAUDE.md`) containing:

- The SUBAGENT CONSTITUTIONAL PREAMBLE block** — a verbatim, self-contained text block covering the top-10 mandatory constraints that cannot be pattern-matched by the hook alone (anti-bluff intent, resource caps, evidence honesty, continuation maintenance, hardcoding prohibition, distribute gate, remote policy, no-guessing vocabulary, etc.) plus the constraint classes that the hook already enforces (emulators, force-push, sudo, host-power) for completeness.
- Clear instruction that the orchestrator MUST paste this block verbatim into every subagent dispatch.**

The preamble is the semantic complement to the hook: the hook catches command-class violations; the preamble pre-informs the agent of all rules it must follow before it even issues a command. Together they eliminate the “agent forgot because the orchestrator didn’t paste it” failure class.

Anti-bluff requirement: the preamble **MUST** be structured so it can be mechanically verified to be present in docs/AGENT_GUARDRAILS.md by check-constitution.sh (or the project-equivalent sweep); the document **MUST** contain the anchor literal 11.4.109 once this clause lands. Paired §1.1 mutation: remove the preamble's hook-reference sentence → constitution checker detects absence → FAIL.

(C) Orchestrator Pre-Action Checklist — guard against self-forgetting.

The same docs/AGENT_GUARDRAILS.md document **MUST** contain an **ORCHESTRATOR PRE-ACTION CHECKLIST** covering at minimum:

- Before any subagent dispatch: confirm the preamble (sub-rule B) is pasted verbatim.
- Before any emulator/device action: confirm the run routes through the Containers submodule CLI, not host-direct; confirm no live device is targeted; confirm the gate host is eligible (otherwise honestly BLOCKED).
- Before any distribute action: confirm Challenge Tests EXECUTED (not compiled) against the exact artifact; version code bumped; CHANGELOG entry present; debug-stage evidence present if two-stage distribute applies.
- Before any push / destructive git action: confirm no force flags without per-operation operator approval; remote is the approved set only; for history rewrite, hardlinked .git backup made first.
- Before any host-affecting command: confirm the command is not in the host-power blocked class.

Anti-bluff requirement: the checklist **MUST** be present in the preamble document and verifiable by check-constitution.sh. Each checklist item is traceable to a governing clause cited inline.

Mechanical enforcement.

Consuming-project check-constitution.sh (or verify-all-constitution-rules.sh) **MUST** verify ALL of:

1. constitution/scripts/hooks/guard-forbidden-commands.sh exists at the canonical path in the constitution submodule.
2. The consumer's AI-agent settings file (.claude/settings.json or equivalent) contains a PreToolUse hook entry referencing the canonical script path.
3. docs/AGENT_GUARDRAILS.md (or the equivalent preamble path) exists, contains the SUBAGENT CONSTITUTIONAL PREAMBLE heading, and contains the ORCHESTRATOR PRE-ACTION CHECKLIST heading.
4. A hermetic test for the hook exists (at minimum, a file under tests/hooks/ or equivalent) and passes.
5. The anchor literal 11.4.109 is present in every per-scope governance file (CLAUDE.md, AGENTS.md, QWEN.md) of the consuming project and its owned submodules per §11.4.35 inheritance.

Pre-build gate CM-ANTI-FORGETTING-ENFORCEMENT enforces (1)–(4). Propagation gate CM-COVENANT-114-109-PROPAGATION enforces (5). Paired §1.1 meta-test mutations: (a) remove the PreToolUse hook entry from the agent settings file → gate (2) FAILs; (b) delete docs/AGENT_GUARDRAILS.md → gate (3) FAILs; (c) remove the canonical hook script from the constitution submodule → gate (1) FAILs; (d) strip the 11.4.109 literal from a consumer CLAUDE.md → propagation gate FAILs.

Inheritance.

Applies recursively to every submodule, every feature, every new artifact, every project consuming this constitution. Submodule constitutions MAY add stricter rules (additional blocked command classes, additional checklist items, stricter test counts) but MUST NOT relax any clause. Per §11.4.35: this universal clause lives in the constitution submodule; consuming projects' CLAUDE.md / AGENTS.md / QWEN.md carry the inheritance pointer-block per §11.4.35 + the 11.4.109 literal per the propagation gate.

The principle this clause enshrines (in the fewest possible words): A constraint that depends on an agent remembering it is not a constraint — it is a hope. The PreToolUse hook converts the most dangerous command classes into hard-blocked failures regardless of agent memory. The preamble document converts the full ruleset into something the orchestrator can inject rather than recite from training. Together they are the minimum viable “anti-forgetting” enforcement posture for any AI-agent-assisted project under this constitution.

Composes with §11.4 + §11.4.1..§11.4.16 (anti-bluff covenant — the hook and preamble are the floor and ceiling of the anti-bluff enforcement layer), §11.4.6 (no-guessing — the escape hatch requires a <reason> exactly because undocumented exceptions are a guessing posture), §11.4.10 (credentials — the host-power class is categorically non-overridable, analogous to credential leaks), §11.4.75 (mechanical enforcement — §11.4.109 is the AI-agent-side specialisation of §11.4.75's git-hook mechanical enforcement; together they cover the full automated-and-agentic surface), §11.4.76 (Containers mandate — the emulator gate in the hook enforces §11.4.76 at the tool-call boundary), §11.4.78 / §11.4.79 / §11.4.80 (CodeGraph — the hook passes codegraph MCP calls through untouched; the preamble includes the anti-forgetting constraint about CodeGraph use), §11.4.81 (cross-platform parity — the hook script is portable bash/awk with no OS-specific primitives; per-OS blocked-class equivalents are addable via consuming-project extension), §11.4.84 (working-tree quiescence — the pre-action checklist instructs agents to grep for mutation markers before staging), §11.4.98 (full-automation mandate — the hook's test suite MUST itself be fully self-driving end-to-end), §11.4.102 (systematic-debugging — any hook violation that surfaces should trigger the systematic-debugging arc, not a quick workaround), §12 (host-session-safety — the host-power class is non-overridable; §12 is the constitutional root of that prohibition).

Classification: universal (§11.4.17) — the PreToolUse guard hook + subagent preamble + orchestrator checklist pattern is vendor-neutral (defined against the JSON stdin contract that every Claude Code hook receives; the pattern ports to any agent runtime that exposes a pre-tool-call interception point) and reusable across ANY project that uses AI coding agents. The specific blocked command-class list at the UNIVERSAL layer covers the four classes that are unconditionally forbidden by this constitution (host-power, privilege escalation, force-push/bypass, raw emulator host-direct); consuming projects extend the list with project-specific additions per §11.4.35.

Non-compliance is a release blocker. No escape hatch — no --skip-pretooluse-hook, --no-guardrails-doc, --anti-forgetting-optional, --single-layer-sufficient flag exists. A project whose agent can forget a constitutional constraint and execute the forbidden command is constitutionally undefended, irrespective of how well the docs are written.

Canonical authority: this Constitution.md §11.4.109 in the HelixConstitution submodule; reference implementation constitution/scripts/hooks/guard-forbidden-commands.sh + reference preamble document constitution/docs/AGENT_GUARDRAILS.md (path within the submodule). Consuming projects inherit by reference per §11.4.80.

§11.4.110 — Pre-build build-readiness verdict + change-impact clash detection mandate (operator mandate, 2026-06-03)

Short tag: pre-build-readiness-verdict.

§-slot history note. Drafted as §11.4.109 in docs/research/prebuild_rigor_20260603/PROPOSED_CONSTITUTION_ANCHOR.md, renumbered to §11.4.110 per §11.4.71 fetch-before-push after the concurrent landing of §11.4.109 (Anti-Forgetting Enforcement, commit 1d9e5d6). Same collision-resolution pattern as §11.4.75 / §11.4.76.

Forensic anchor (genericised, 2026-06-03). A fix shipped a new system-property *read* in a code/init change but introduced no matching security-policy grant and no property-context type entry for it — so the read was *silently denied* at runtime and the feature was dead, while every pre-build gate stayed green because the gate only grepped the source file for the change. The defect class generalises: a change introduces a new dependency on a *second* artifact (security policy, context file, service registry, interface freeze-snapshot, symbol table, build-graph node) that the pre-build never cross-checks, so the change is “ready” by every existing gate yet broken the instant it runs. This is the SOURCE→ARTIFACT half of §11.4.108’s four-layer gap, shifted *left* to pre-build time: most of these clashes are **statically catchable from the change diff itself, before any build**.

The architectural flaw (FACT, traced). Pre-build gates that validate each artifact *in isolation* (the policy file is internally consistent; the freeze-snapshot is well-formed) do NOT cross-check that a change introducing a new dependency *also* introduced the matching second artifact. The cross-check is **diff-driven** (its input is the change set, not any single file) and is therefore a capability no per-file gate provides. Without it, the only protection is author-vigilance (“remember to also add the policy line”), which §11.4.75 already established is the failure mode that mechanical enforcement exists to replace.

The mandate (ALL must hold).

1. **A single deterministic READY-FOR-BUILD verdict is mandatory.** Before any artifact rebuild, the project MUST emit ONE machine-checkable verdict — *the codebase IS / IS NOT ready for build* — backed by a captured-evidence bundle enumerating every gate family’s status. The rebuild orchestrator MUST refuse to start the build unless the verdict is READY (non-zero exit on NOT-READY). A green individual gate is necessary but not the verdict; the verdict is the aggregate function and cannot report READY while any family reports FAIL.
2. **A diff-driven change-impact + conflict/clash detector is mandatory.** For every change in the batch, the pre-build MUST cross-check that any newly-introduced dependency on a *second* artifact is satisfied — at minimum: a new system-property read $\Leftrightarrow$ the property is typed in the platform’s property-context

registry AND the reading domain has the matching read-grant; a new registered service $\Leftrightarrow$ its service-context entry exists; a new declared startup/init service $\Leftrightarrow$ its security label is present; a new/changed stable interface $\Leftrightarrow$ its freeze-snapshot/version was updated in the same change; a new native-library dependency $\Leftrightarrow$ the named library resolves to a known module or prebuilt; a new security-policy rule $\Leftrightarrow$ every type/attribute it uses is defined; and two changes in the batch touching the SAME function / resource / configuration seam $\Leftrightarrow$ the collision is explicitly acknowledged (the cross-fix-interaction class). Each sub-detector is independently anti-bluff (clause 5).

3. **Coverage-completeness is a gate, not a convention.** Every changed source file in the batch MUST map to ≥ 1 pre-build gate + ≥ 1 deployed-target test + ≥ 1 paired §1.1 mutation (the four-layer §11.4.4(b) floor). A baseline file records the legacy tail of known holes; the gate FAILs on any NEW hole and the covered-ratio ratchets upward over phases per §11.4.50. A change with no owning gate/test/mutation is a coverage hole and blocks the verdict.
4. **Two-speed honesty — grep-speed vs REQUIRES_BUILD — is mandatory.** The regime MUST separate always-on grep-speed gates (run every commit, seconds) from REQUIRES_BUILD heavy gates (the build-graph parse-only dry-run; full security-policy neverallow compilation; built-symbol ABI diff). Heavy gates run as diff-gated opt-in stages (triggered only when the change touches build-graph / policy / interface files) and MUST run bounded per the host-session-safety budget (§12.6/§12.7) or inside the project's containerized build. Conflating the two — making the grep path slow, or claiming grep proves graph/neverallow correctness — is forbidden per §11.4.6.
5. **Every gate and analyzer is anti-bluff by paired mutation.** Each gate family AND each wired third-party analyzer (its exit-code-on-failure semantics included) MUST ship a paired §1.1 mutation that plants the exact defect class and asserts the gate FAILs (e.g. a property read with no policy $\rightarrow$ the clash detector FAILs; a syntactically-broken build-graph fragment $\rightarrow$ the dry-run reports FAIL; an analyzer that exits 0 on a planted violation is itself the bluff the mutation catches). An analyzer wired without its mutation pair is a §11.4 bluff surface.
6. **Honest boundary — pre-build catches statically-catchable defects only.** A READY verdict proves the change is internally consistent and ready to build; it does NOT prove the feature works on the deployed target. The correct claim is: *after the pre-build passes, the residual defects deployed-target / manual testing can find are runtime/hardware defects, NOT statically-catchable (preventable) ones* — the regime makes the **preventable** class empty, not the **all-defects** class. Overclaiming that manual testing will find nothing is itself a §11.4.6 violation. Build-graph correctness, full neverallow safety, and ABI correctness are proven ONLY by the REQUIRES_BUILD families (clause 4); the grep tier proves harness-presence + symbol-definedness, never full correctness. Runtime/USER-VISIBLE correctness remains entirely §11.4.108's four-layer job.

Classification: universal (§11.4.17) — the single ready-for-build verdict, the diff-driven change-impact/clash detector, the coverage-completeness gate, the grep-speed-vs-REQUIRES_BUILD two-speed split, the per-gate paired mutation, and the statically-catchable-only honest boundary are platform-neutral pre-build-rigor disciplines reusable by ANY project that builds an artifact from source; the

consuming project supplies its concrete property/service/policy registries, its build-graph parse-only command, its interface-freeze mechanism, and its changed-file-to-gate mapping per §11.4.35.

Composes with §11.4.1 (FAIL-bluffs forbidden — a clash-detector that crashes on a malformed diff and FAILs is as misleading as one that passes silently), §11.4.4 (test-interrupt-on-discovery — a NOT-READY verdict triggers STOP-and-fix before the build), §11.4.6 (no-guessing — the two-speed honest boundary; the grep tier MUST NOT claim correctness it cannot prove), §11.4.9 (batch-source-fixes-before-rebuild — the verdict gates exactly that rebuild), §11.4.27 (no-fakes-beyond-unit — wired analyzers exercise the real change), §11.4.50 (deterministic consistency — the coverage-completeness ratchet shares §11.4.50's threshold mechanism and N-iteration determinism), §11.4.67 (target-shell-parseability — the clash-detector helper itself parses under its target shell), §11.4.75 (mechanical enforcement — §11.4.110 is the mechanical replacement for author-vigilance on the clash class), §11.4.92 (multi-pass change-evaluation — the verdict is the mechanical companion to Pass 1/2/3), §11.4.108 (four-layer fix-verification — §11.4.110 is the SOURCE→ARTIFACT half shifted left to pre-build time; §11.4.108 owns the RUNTIME-ON-CLEAN-TARGET→USER-VISIBLE half; together they span all four layers).

Propagation. Propagation gate CM-COVENANT-114-110-PROPAGATION enforces the literal anchor 11.4.110 across the consumer fleet (parent + owned-submodule CLAUDE.md / AGENTS.md / QWEN.md); paired §1.1 meta-test mutation strips the 11.4.110 literal from a consumer file → the gate FAILs. Recommended per-family gates CM-READY-FOR-BUILD-VERDICT (single deterministic verdict gates the rebuild), CM-CHANGE-IMPACT-CLASH-DETECTOR (diff-driven second-artifact cross-check), CM-COVERAGE-COMPLETENESS-GATE (changed-file $\rightleftharpoons$ gate+test+mutation with baseline ratchet), CM-BUILDGRAPH-DRYRUN-WIRED + CM-SEPOLICY-NEVERALLOW-WIRED (REQUIRES_BUILD opt-in stages exist + bounded-safe); each with a paired §1.1 mutation per clause 5. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md §11.4.110 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no --source-green-is-ready, --skip-clash-detector, --skip-coverage-gate, --no-ready-verdict, --grep-proves-neverallow, --skip-buildgraph-dryrun, --build-without-ready-verdict flag exists.

§11.4.111 — Resolve-by-stable-name-not-by-enumeration-index mandate (research-derived, 2026-06-03)

Short tag: resolve-by-stable-name.

Forensic anchor (genericised, 2026-06-03). A platform bound an audio output to a kernel-enumerated device index (card=0). A second device of a different class enumerated FIRST at boot and took slot 0, shifting the intended device to slot 1. The static index binding now pointed at the wrong device: the policy layer failed to attach the intended sink, mis-assigned its TYPE, and the user-facing output switcher labelled the AV receiver “Wired Headphone” while routing collapsed to stereo. The lower layer of the SAME stack already resolved the device correctly **by name** (scanning /proc/asound/card*/id for the controller name) — proving the

brittleness was the *index* binding, not the resolution capability. The defect class generalises to every enumerated resource whose ordinal is assigned at discovery/ boot/hotplug time and is therefore non-deterministic across reboots, device additions, and topology changes.

The mandate. Any binding between a configuration / policy / code site and a hardware device, resource handle, or enumerated entity — audio cards, display connectors, network interfaces, storage devices, GPU render nodes, input devices, camera indices, container/process slots, any registry whose members are assigned an ordinal at discovery time — MUST resolve the target by a **stable identifier** (name, UUID, serial, label, controller-name, content-derived hash, sink-reported identity) and MUST NOT resolve it by an enumeration index / ordinal / slot number, UNLESS the platform documents that ordinal as deterministically pinned (and the pinning mechanism is itself part of the binding, captured and asserted). Where a stable identifier exists at one layer of a stack, every other layer binding the same resource MUST use the same identifier — mixed by-name-here / by-index-there is the structural weak link this anchor forbids. Where the platform offers a sink-reported / capability-derived identity (e.g. an EDID/ELD-style capability descriptor), preferring it makes the binding robust to the broadest class of topology change.

Honest boundary. When the platform genuinely provides no stable identifier and only an ordinal exists, the binding MUST (a) pin the ordinal deterministically via the platform's own pinning mechanism, (b) capture that pin as part of the binding's runtime signature per §11.4.108, and (c) document the residual fragility as UNCONFIRMED: -class risk per §11.4.6 — never silently trust an unpinned ordinal.

Classification: universal (§11.4.17) — index-vs-name resource resolution is a platform-neutral binding-robustness discipline reusable by ANY project that binds to enumerated hardware / resources / handles; the consuming project supplies its concrete stable-identifier mechanism (ALSA card name, DRM connector name, NIC predictable name, block-device UUID, etc.) and the specific layers that must agree per §11.4.35.

Composes with §11.4.6 (no-guessing — “the index is *usually* stable” is the exact guess this anchor forbids; prove the pin or bind by name), §11.4.8 (deep-web-research — mature platform stacks resolve by name/UUID; reproducing a known-brittle index binding when the by-name path is documented is a §11.4.8 violation by omission), §11.4.69 (sink-side positive-evidence — the by-name binding's correctness is verified by the downstream/sink-reported identity, not by the local config), §11.4.108 (four-layer fix-verification — the resolved-by-name binding's runtime signature is asserted on a clean target across the topology that originally broke the index binding), §11.4.110 (pre-build clash detection — a new ordinal-based binding in a diff is a statically-catchable clash class the change-impact detector flags).

Propagation. Propagation gate CM-COVENANT-114-111-PROPAGATION enforces the literal anchor 11.4.111 across the consumer fleet; paired §1.1 meta-test mutation strips the literal → the gate FAILs. Recommended per-family gate CM-RESOLVE-BY-NAME-NOT-INDEX (diff-driven scan flagging a new enumeration-index binding to a resource that exposes a stable identifier, unless an accompanying deterministic-pin + runtime-signature is present); paired §1.1 mutation plants an unpinned index binding → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md §11.4.111 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no --allow-index-binding, --ordinal-is-stable-enough, --skip-name-resolution, --trust-unpinned-index flag exists.

§11.4.112 — Structural-impossibility won't-fix classification mandate (research-derived, 2026-06-03)

Short tag: structural-impossibility-wont-fix.

Forensic anchor (genericised, 2026-06-03). Deep research (§11.4.8) into relocating protected (content-protection / secure-surface) video to a secondary display proved — from authoritative platform/HDCP documentation plus reproducible captured behaviour (a secure surface is blanked on any output lacking the secure flag; mirror/screenshot of a secure layer returns black) — that the goal is **structurally impossible by platform design**, not a missing feature or an unsolved engineering problem. Without a durable classification, such a goal is re-investigated every cycle: each new agent / operator re-reads the same sources, re-runs the same probes, re-derives the same impossibility, and burns the same effort — an anti-pattern that compounds across cycles.

The mandate. When deep research per §11.4.8 PROVES (with cited authoritative sources AND, where applicable, reproducible captured evidence) that a goal is structurally impossible on the target platform — i.e. forbidden by the platform's design / a hardware/protocol constraint / a documented kernel-or-API limitation, not merely unimplemented or hard — the goal MUST be: (1) **classified Won't-fix and closed** per the §11.4.90 terminal-status discipline, with obsolescence/closure reason structurally-impossible (a value in the §11.4.90 closed reason vocabulary); (2) **documented with the impossibility evidence** — the cited authoritative source URLs (per §11.4.99 latest-source verification where the platform docs may evolve) PLUS the reproducible probe/captured evidence that demonstrates the constraint, recorded in the tracker entry and the relevant docs/ guide so the verdict is auditable; (3) **NOT re-attempted in future cycles** — the closed entry is the canonical answer; a future agent that re-opens it MUST cite NEW evidence that the platform constraint changed (per §11.4.34 reopened-source attribution + §11.4.7 demotion-evidence), never merely re-derive the same impossibility; (4) **paired with the correct posture** — the entry MUST state what the project DOES instead (the supported-within-the-constraint behaviour), so “impossible” is never confused with “broken / unhandled”.

Honest boundary (§11.4.6). Won't-fix: structurally-impossible is reserved for *proven* platform/hardware/protocol impossibility. “We could not find a way” / “it is very hard” / “no time” are NOT structural impossibility — they are Operator-blocked (§11.4.21) or open work, and mislabelling them won't-fix to avoid the work is a §11.4 bluff at the planning layer. A future platform change CAN make the impossible possible; the classification is durable but not eternal, and reopening on NEW cited evidence is correct.

Classification: universal (§11.4.17) — proving-then-durably-classifying structural impossibility is a platform-neutral effort-conservation + honesty discipline reusable by ANY project; the consuming project supplies the specific impossible goal, its platform constraint, and the cited evidence per §11.4.35.

Composes with §11.4.6 (no-guessing — the impossibility is stated as FACT with cited evidence, never as “probably can’t”), §11.4.7 (demotion-evidence — reopening a won’t-fix requires NEW positive evidence the constraint changed), §11.4.8 (deep-web-research — the proof is the research output; “NO external solution found — structurally impossible” is the §11.4.8 citation), §11.4.34 (reopened-source attribution — a reopen of a structurally-impossible item attributes who reopened + the new evidence), §11.4.90 (Obsolete/terminal-status — structurally-impossible is a closure reason in the §11.4.90 closed vocabulary; the colorizer marks it like any closed item), §11.4.99 (latest-source verification — the cited platform docs MUST be the latest where they can evolve, so the impossibility verdict is not stale).

Propagation. Propagation gate CM-COVENANT-114-112-PROPAGATION enforces the literal anchor 11.4.112 across the consumer fleet; paired §1.1 meta-test mutation strips the literal → the gate FAILs. Recommended per-family gate CM-WONT-FIX-STRUCTURAL-IMPOSSIBILITY (every tracker entry with closure reason structurally-impossible carries cited authoritative-source evidence + a stated correct-posture line); paired §1.1 mutation strips the evidence citation from such an entry → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md §11.4.112 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no --wont-fix-without-proof, --reattempt-closed-impossible, --skip-impossibility-evidence, --impossible-equals-broken flag exists.

§11.4.113 — Absolute no-force-push + merge-onto-latest-main mandate (User mandate, 2026-06-03)

Short tag: absolute-no-force-push.

Forensic anchor — verbatim user mandate (2026-06-03):

“Any force-push is strictly forbidden! We must for every Submodule take as a base latest commit on Submodule’s main (or master) branch, then on top of it carefully to merge all changes that have to be pushed! Once all merging is carefully done we perform commit and push to all Submodule’s upstreams!”

The mandate — force-push is STRICTLY FORBIDDEN, with NO exception. Any force-push — `git push --force`, `git push --force-with-lease`, `git push +<ref>`, or any history-rewriting operation that overwrites a remote ref non-fast-forward — is forbidden against every repository this Constitution governs (every consuming project’s main repo, this constitution submodule, every owned submodule, every nested submodule, every upstream of each). There is no escape hatch, no operator-approval path, no “after a merge-first audit” path. The merge-onto-latest-main procedure below is ALWAYS available, so a force-push is NEVER necessary; eliminating the need eliminates the operation.

The mandated 6-step integration procedure (for every repo/submodule whose local has commits to publish OR whose mirrors have diverged):

1. **Fetch all remotes** — `git fetch --all --prune --tags` against origin + every configured upstream; capture the output as evidence.
2. **Set the base to the LATEST commit on the canonical main/master branch** — the most-advanced mirror tip across all upstreams. The integration target is that tip, never a stale local branch point.
3. **Carefully MERGE every change that must be published on top of that base** — a union merge that preserves BOTH sides (local commits AND every remote-side commit). NEVER `-s ours` (it silently discards the remote side), NEVER a rebase/reset/—allow-unrelated-histories that could drop commits (per §9 no-commit-loss). The merge commit MUST descend from every mirror tip so that every subsequent push is a clean fast-forward.
4. **Resolve every conflict carefully** so nothing is broken, lost, or corrupted on either side — no conflict markers anywhere (`grep -rn '^<<<<<< \|^===== $\|^>>>>>> '` returns empty across governance + source + test files), no file silently dropped, every previously-passing gate/test still passes on the integrated tree (per §11.4.4 + §11.4.40), every captured-evidence artefact still validates.
5. **Commit the merge** — stage only the intended files (NEVER `git add -A` inside a submodule with untracked build content per §11.4.30), commit the carefully-resolved merge.
6. **Push the result to ALL upstreams** — each push is a fast-forward because the merge commit descends from every mirror tip, so NO force is ever needed. If an upstream still reports non-fast-forward (a sibling landed work between step 1 and the push), do NOT force — return to step 1 for that upstream, fetch its new tip, MERGE it (union, preserve both), re-validate step 4, and re-push.

Tightens §11.4.41 / §11.4.71 / §9.2 / CONST-043 — force-push escape hatch REMOVED. §11.4.41 (Pre-Force-Push Merge-First Mandate), §11.4.71 (Pre-Push Fetch + Investigate + Integrate), §9.2 (Force-push requires explicit user authorization), and CONST-043 previously PERMITTED a force-push after a merge-first pipeline + per-operation operator approval. §11.4.113 supersedes that permitting stance: even WITH operator approval, even AFTER a clean merge-first audit, a force-push is forbidden — because the merge-onto-latest-main path in this clause is always available and always sufficient, force is never necessary and therefore never authorised. Those clauses' merge-first/fetch-first machinery REMAINS in force as the integration discipline; only their terminal "...then force-push" step is struck. §9.2's destructive-operation §9 backup discipline still applies to any non-push history operation, but a history-rewriting *push* is simply not performed.

No escape hatch. No `--force`, no `--force-with-lease`, no `+<ref> push`, no `--no-verify` to slip a forced ref past a hook, no "operator authorised the force-push" path, no "lease semantics unavailable so plain `-force` is OK" path. A force-push that lands on any remote is a critical defect regardless of how briefly it persisted, severity-equivalent to a §11.4 PASS-bluff at the data-safety layer (per §9).

Classification: universal (§11.4.17) — an absolute-no-force-push + merge-onto-latest-main integration discipline is platform-neutral and reusable across every repository in every consuming project; nothing here references a particular project, vendor, hardware, or region.

Composes with §2.1 (multi-upstream push is the norm — step 6 fans out to all), §9 / §9.2 (absolute data safety — §11.4.113 is the no-loss push discipline that makes force-push unnecessary), §11.4.4 (test-interrupt + clean-baseline retest — step 4 audit), §11.4.6 (no-guessing — remote state is not knowable without the step-1 fetch), §11.4.26 (constitution-submodule update workflow — its conflict-resolution step is this procedure), §11.4.37 (fetch-before-edit — step 1 generalised to fetch-before-push), §11.4.40 (full-suite retest authority for the integrated tree), §11.4.41 (force-push merge-first — TIGHTENED: merge-first stays, the force-push step is removed), §11.4.71 (pre-push fetch + integrate — TIGHTENED the same way), §11.4.88 (background-push — the detached push still follows this fast-forward-only discipline), CONST-043 (per-operation authorization — TIGHTENED: no force-push is authorisable).

Propagation. Propagation gate CM-COVENANT-114-113-PROPAGATION enforces the literal anchor 11.4.113 across the consumer fleet; paired §1.1 meta-test mutation strips the literal → the gate FAILs. Recommended per-family gate CM-NO-FORCE-PUSH-ABSOLUTE (scans tracked scripts + hooks for any push --force / push --force-with-lease / push +<ref> invocation and rejects it; a §11.4.109-class PreToolUse guard blocks the force-push command class at the tool-call boundary); paired §1.1 mutation injects a git push --force into a tracked script → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md §11.4.113 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no --force, --force-with-lease, --allow-force-push, --force-push-authorised, --skip-merge-onto-main flag exists.

§11.4.114 — Last-known-good-tag regression isolation mandate (1.1.8-dev remediation, 2026-06-03)

Short tag: last-known-good-tag-isolation.

Forensic anchor (genericised, 2026-06-03). A user reported that a previously-working feature (multichannel HDMI audio) was broken on the current build. The operator's lead — “the feature WORKS and passes all tests at the last release tag” — converted an open-ended root-cause hunt into a targeted regression repair: diffing the broken working tree against the last-known-good tag's commit isolated EVERY reported audio defect (multichannel downmix + output-device label mismapping) to ONE post-tag burndown batch (AudioDeviceInventory.connectDualHdmiPorts_1 dual-HDMI-connect) in minutes, not hours. The tagged-good version is the oracle: behaviour that works there and is broken now is, by definition, a regression introduced in the diff between the two.

The mandate. When a previously-working feature, behaviour, or capability is observed broken, the FIRST diagnostic action MUST be to identify the last release tag (or last commit / build / deployment) at which that feature was KNOWN-GOOD, and diff/bisect the broken state against it — BEFORE any open-ended root-cause investigation, BEFORE any speculative fix. The known-good revision is the regression oracle: (1) it bounds the search space to the commits between known-

good and now (a `git diff <good-tag>..HEAD --stat` of the feature's files is the captured evidence); (2) it tells you the fix is a regression REPAIR (restore/forward-fix the broken sub-part), not a from-scratch design problem; (3) each suspect file is compared against its known-good version as the behavioural oracle. When the operator volunteers a known-good tag, that lead is load-bearing and MUST be acted on first. Where the bisect surface is large, `git bisect` against the feature's autonomous RED test (per §11.4.115) mechanises the search.

Forward-fix vs wholesale-revert. Once the regressing change is isolated, the default is a SURGICAL forward-fix — keep the new features the post-good-tag work introduced, revert ONLY the broken sub-part where the known-good version was correct — UNLESS the operator prefers wholesale revert. A wholesale revert of a multi-feature batch to reach the known-good state is a last resort that loses the batch's other (working) features; prefer isolating and repairing the single regressing seam. Every forward-fix MUST prove (via the §11.4.115 RED→GREEN test) that the feature now matches the known-good behaviour.

Honest boundary (§11.4.6). “It worked before” is a HYPOTHESIS until the known-good tag is identified AND the feature is confirmed working there (re-run the feature's test against the known-good build, or cite the captured-evidence from when that tag shipped). A feature that NEVER worked is not a regression and this anchor's oracle does not apply — investigate as new work. “Probably regressed in the last batch” without the diff is a §11.4.6 guess.

Classification: universal (§11.4.17) — diff-against-last-known-good is a platform-neutral regression-isolation discipline reusable by ANY versioned project with release tags / commit history; the consuming project supplies its tag-naming convention and per-feature known-good oracle per §11.4.35.

Composes with §11.4.4 (test-interrupt-on-discovery — the isolation runs as the first step of the systematic-debugging Phase 1 the STOP triggers), §11.4.6 (no-guessing — the diff is the FACT-grade evidence that replaces “probably the last batch”), §11.4.7 (demotion-evidence — confirming known-good requires positive evidence at that tag), §11.4.40 (full-suite retest authority — the known-good tag is the baseline a release-gate retest restores toward), §11.4.43 (TDD-fix — the §11.4.115 RED test reproduces the regression the diff isolated), §11.4.102 (systematic-debugging — last-known-good isolation IS the Phase 1 root-cause technique for regressions specifically), §11.4.108 (four-layer fix-verification — the forward-fix's runtime signature is verified against the known-good behaviour on a clean target).

Propagation. Propagation gate CM-COVENANT-114-114-PROPAGATION enforces the literal anchor 11.4.114 across the consumer fleet; paired §1.1 meta-test mutation strips the literal → the gate FAILs. Recommended per-family gate CM-REGRESSION-ISOLATED-AGAINST-KNOWN-GOOD (every tracker entry classified as a regression cites the last-known-good tag + the diff range it was isolated to); paired §1.1 mutation strips the known-good citation from a regression entry → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md §11.4.114 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--skip-known-good-diff`, `--root-cause-from-scratch`, `--assume-regression-source`, `--wholesale-revert-without-isolation` flag exists.

§11.4.115 — RED-baseline-on-the-broken-artifact + polarity-switch mandate (1.1.8-dev remediation, 2026-06-03)

Short tag: red-baseline-polarity-switch.

Forensic anchor (genericised, 2026-06-03). Per §11.4.43 a fix opens with a RED test. The 1.1.8-dev remediation sharpened HOW: the autonomous RED test was authored to REPRODUCE each live defect on the CURRENT pre-fix build — proving the defect is real, live, and captured (video flips to primary on reconfigure 3/3; audio 5.1 → Arvus PCM2/0 ×3; subtitle absent on secondary) — with a single environment flag (RED_MODE=1 default) that, set to 0 post-fix, flips the SAME test to GREEN-guard polarity. One test source, two roles: the reproduction-of-the-defect (RED) and the permanent regression-guard (GREEN). No separate happy-path test is authored; the test that catches the bug IS the test that guards against its return.

The mandate. Every §11.4.43 RED test MUST be authored to reproduce the defect on the CURRENT, pre-fix artifact (the actual broken build / deployment), capturing positive evidence per §11.4.5 / §11.4.69 / §11.4.107 that the defect is genuinely present — never a synthetic / hypothetical failure the fix is then written to agree with. The SAME test source MUST carry a single polarity switch (an env flag / parameter, canonical RED_MODE, default 1 = reproduce-and-assert-defect-present) that, flipped to 0 post-fix, converts the test into the GREEN regression-guard asserting the defect is ABSENT. The RED run (against the broken artifact) and the GREEN run (against the fixed artifact, per §11.4.108 on a clean target) MUST both be captured as evidence: RED-then-GREEN is the proof the test catches the bug AND the fix closes it. A separate happy-path test that never failed on the broken artifact is forbidden as the primary guard — it demonstrates only that the test agrees with the fix (the §11.4.43 PASS-bluff), not that it catches the defect.

Why one test, two roles. A test authored only as a post-fix happy-path can pass on a build where the feature is silently broken (it asserts the fix's intended state, not the defect's absence). A RED-first test that was OBSERVED failing on the broken artifact, then flipped to GREEN-guard, is provably sensitive to the exact defect — its RED capture is the falsification record. The polarity switch keeps the two roles in ONE source so they cannot drift apart (a separate happy-path test can be edited to stop testing the real condition; the RED-evidence binds the GREEN guard to the actual defect).

Honest boundary (§11.4.6). If the defect cannot be reproduced on the current artifact (the RED run does not fail), that is a finding, not a license to write a GREEN-only test: either the defect is not present on this build (close per §11.4.7 with the negative evidence) or the test does not exercise the real condition (fix the test). A RED test that “passes” on the known-broken artifact is itself a §11.4.1 FAIL-bluff inverse — it proves the test is blind, not that the feature works.

Classification: universal (§11.4.17) — RED-baseline-on-the-broken-artifact + a single-source polarity switch is a platform-neutral test-construction discipline reusable by ANY project doing TDD-fix; the consuming project supplies its polarity-flag mechanism (env var, build flag, parameter) and per-defect reproduction driver per §11.4.35. Strict refinement of §11.4.43 step RED.

Composes with §11.4.1 (FAIL-bluffs — a RED test that passes on the broken artifact is blind), §11.4.2 / §11.4.5 / §11.4.69 / §11.4.107 (the RED reproduction captures positive defect-present evidence), §11.4.4 (the RED test is the test-interrupt reproduction), §11.4.7 (RED-not-reproducible is a demotion that needs negative evidence), §11.4.43 (refines its RED step — RED is on the real broken artifact, GREEN is the same test flipped), §11.4.50 (deterministic consistency — both RED and GREEN runs are N-iteration stable), §11.4.108 (the GREEN run is the runtime-signature verification on a clean target), §11.4.114 (the RED test mechanises the §11.4.114 bisection against the known-good oracle).

Propagation. Propagation gate CM-COVENANT-114-115-PROPAGATION enforces the literal anchor 11.4.115 across the consumer fleet; paired §1.1 meta-test mutation strips the literal → the gate FAILs. Recommended per-family gate CM-RED-POLARITY-SWITCH-PRESENT (every regression-fix RED test carries a single polarity switch + a captured RED-on-broken-artifact run before the fix); paired §1.1 mutation removes the polarity switch (splitting RED and GREEN into a happy-path-only guard) → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md §11.4.115 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no --green-only-test, --skip-red-reproduction, --separate-happy-path-suffices, --synthetic-red-OK flag exists.

§11.4.116 — Real-time conductor[?] autonomous-test-framework sync channel mandate (1.1.8-dev remediation, 2026-06-03)

Short tag: conductor-test-sync-channel.

Forensic anchor (genericised, 2026-06-03). An autonomous test / QA framework (HelixQA) was wired into the conductor's loop. For the conductor (the orchestrating agent / operator) to stay in real-time sync with a long-running autonomous session — knowing which phase / challenge is executing, what evidence was captured, which LLM/vision call ran, and the per-item PASS/FAIL/SKIP/OPERATOR-BLOCKED verdict the moment it lands — the framework exposed a structured append-only JSONL event stream (conduit.events.jsonl) plus an atomically-rewritten status snapshot (conduit.status.json) that the conductor tails live. Without it, the conductor either polls blindly (idle waiting, §11.4.94 violation) or learns the outcome only at session end (no mid-session course-correction, no §11.4.4 test-interrupt-on-discovery).

The mandate. Any autonomous, long-running test / QA / validation framework that an external orchestrator (conductor agent or operator) depends on for real-time decisions MUST expose a real-time sync channel with two parts: (1) a **structured append-only event stream** (JSONL or equivalent, one event per line, never rewritten) emitting at minimum session-start / phase-transition / per-test-or-challenge-start / captured-evidence-path / external-call (LLM / vision / sink-probe) / error / per-item-verdict events; (2) an **atomically-rewritten status snapshot** (a single small JSON/file written via write-temp-then-rename so a reader never observes a torn write) carrying the current session/phase/item + running counters + last verdict. Verdicts MUST use the closed vocabulary PASS / FAIL / SKIP /

OPERATOR-BLOCKED (per §11.4.45 status vocabulary). The conductor tails the stream / snapshot live (`tail -f` / a dedicated monitor) so it stays in real-time sync, can §11.4.4-interrupt on a fresh defect, and never idles blindly per §11.4.94 / §11.4.97. The channel is the framework's contract with the orchestrator, not an afterthought log.

Anti-bluff for the channel itself (§11.4 / §11.4.69). The event stream + snapshot are themselves captured-evidence artefacts: a verdict event **MUST** carry the evidence path that backs it (per §11.4.69 `ab_pass_with_evidence`), so a PASS event with no evidence path is a §11.4 PASS-bluff at the channel layer. A status snapshot that reports PASS while the event stream shows no captured-evidence event for that item is a contradiction the conductor **MUST** treat as FAIL. The channel never invents progress: an item with no start-event cannot have a verdict-event.

Decoupling (§11.4.28). When the framework is an owned, project-agnostic submodule, the channel **MUST** stay project-neutral — the consuming project registers its own data (endpoints, app/package identifiers, sink hosts) at runtime via the framework's public API, never hardcoded into the framework. The channel format is generic; the payloads are consumer-supplied.

Classification: universal (§11.4.17) — a real-time append-only event stream + atomically-rewritten status snapshot is a platform-neutral orchestrator-sync discipline reusable by ANY project whose conductor depends on a long-running autonomous test/QA framework; the consuming project supplies the concrete channel paths and payload schema per §11.4.35.

Composes with §11.4.4 (test-interrupt-on-discovery — the live stream is what lets the conductor STOP the instant a fresh defect's FAIL event lands), §11.4.5 / §11.4.69 (every verdict event cites its captured-evidence path), §11.4.27 (no-fakes — the framework exercises the real system, the channel reports real verdicts), §11.4.28 (owned-submodule decoupling — channel + payloads stay project-agnostic), §11.4.45 (status vocabulary + sink-side evidence reuse), §11.4.52 (autonomous-validation — the channel is how an autonomous session reports without a human watching the screen), §11.4.89 (background test execution — the channel is the conductor's window into a backgrounded long test), §11.4.94 / §11.4.97 (zero-idle — tailing the live channel replaces blind polling/idle waiting).

Propagation. Propagation gate CM-COVENANT-114-116-PROPAGATION enforces the literal anchor 11.4.116 across the consumer fleet; paired §1.1 meta-test mutation strips the literal → the gate FAILs. Recommended per-family gate CM-AUTONOMOUS-FRAMEWORK-SYNC-CHANNEL (the framework emits an append-only event stream + an atomically-rewritten status snapshot, and every verdict event carries an evidence path); paired §1.1 mutation strips the evidence-path field from a verdict event → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md §11.4.116 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--no-sync-channel`, `--end-of-session-report-suffices`, `--verdict-without-evidence-path`, `--torn-status-write-OK` flag exists.

§11.4.117 — Computer-vision / OCR pixel-oracle fallback for non-introspectable UIs mandate (1.1.8-dev remediation, 2026-06-03)

Short tag: cv-ocr-pixel-oracle-fallback.

Forensic anchor (genericised, 2026-06-03). Several apps under test (TV-Compose / leanback streaming apps, canvas-rendered players, games) return a near-empty accessibility / semantic tree — `uiautomator dump` yields a hierarchy with zero clickable nodes, so a hierarchy-driven test cannot find the play tile, cannot read the on-screen caption, cannot confirm what the user sees. The §11.4.52 rule already SKIPS UI-driven when `uiautomator dump | grep -c clickable=true == 0`. The remediation’s lesson: SKIP is not the only answer — when the semantic tree is blank, the test **MUST** fall back to driving + asserting via the PIXELS (computer-vision template-match to locate and tap controls; ROI OCR to read on-screen text), because the pixels are what the end user actually experiences, and a hierarchy-only tool is **NOT** a content oracle.

The mandate. Any test that needs to drive a UI control OR assert on-screen content **MUST NOT** assume the accessibility / semantic / DOM hierarchy is the source of truth for what the user sees. When the hierarchy is blank, partial, or known-unreliable for the app under test (TV-Compose, leanback, canvas/SurfaceView/GL, games, custom-rendered UIs), the test **MUST** fall back to a **pixel oracle**: (1) **drive** input by computer-vision template-match (locate a control by its rendered appearance, then tap its screen coordinates) rather than by a hierarchy node id; (2) **assert** on-screen content by ROI OCR (read the actual rendered text in the region the user reads — caption strip, title, error overlay) with a per-word confidence floor + region-of-interest per §11.4.107(12), never by a hierarchy text attribute that may be absent or stale. The test tool **MUST** be chosen for its ability to **BOTH** drive input AND read pixels — a hierarchy-only tool cannot serve as a content oracle for these UIs. The pixel path is the §11.4.52 “near-empty hierarchy proves UI-driven INFEASIBLE” fallback made constructive: not SKIP, but pixel-drive.

Anti-bluff for the pixel oracle (§11.4.107). The CV/OCR analyzer is itself a potential bluff surface and **MUST** be self-validated per §11.4.107(10): a golden-good fixture (the control present / the expected caption rendered) **MUST PASS** and a golden-bad fixture (control absent / wrong-or-empty caption) **MUST FAIL**, wired into the meta-test. Thresholds (match confidence, OCR confidence floor) **MUST** be calibrated on the project’s own captured frames, not hardcoded from literature (§11.4.107(13) / §11.4.6). An OCR PASS that cannot also FAIL its golden-bad fixture is a §11.4 bluff gate.

Honest boundary (§11.4.6). When **BOTH** the hierarchy is blank AND the pixel oracle is infeasible (e.g. the target output is a protected/secure surface that captures black per §11.4.112, or the content is geo-unreachable), the test SKIPS-with-reason per §11.4.3 (`topology_unsupported` / `operator_attended` / `geo_restricted`), never a fake PASS, and the operator-attended path is a tracked migration item per §11.4.52, not the permanent answer.

Classification: universal (§11.4.17) — driving + asserting via pixels when the semantic tree is blank is a platform-neutral test-oracle discipline reusable by ANY project testing custom-rendered / canvas / game / Compose / leanback UIs; the consuming project supplies its concrete CV/OCR toolchain (template-match library, OCR engine), capture mechanism, and calibrated thresholds per §11.4.35.

Composes with §11.4.5 (captured-evidence quality — OCR overlay census), §11.4.6 (no-guessing — thresholds calibrated, not assumed), §11.4.48 (UI-driven — pixel-drive is the fallback when the hierarchy path is infeasible), §11.4.49 (dual-approach — pixel-drive UI variant complements the intent variant), §11.4.52 (autonomous-validation — strict refinement: near-empty hierarchy → pixel-drive, not only SKIP), §11.4.107 (AV/test-validation — the OCR/CV analyzer is self-validated golden-good/golden-bad, ROI + confidence floor; §11.4.117 is its UI-driving companion), §11.4.112 (structural-impossibility — secure-surface black-capture is the honest boundary where even pixel-drive cannot read content).

Propagation. Propagation gate CM-COVENANT-114-117-PROPAGATION enforces the literal anchor 11.4.117 across the consumer fleet; paired §1.1 meta-test mutation strips the literal → the gate FAILs. Recommended per-family gate CM-CV-OCR-PIXEL-ORACLE-FALLBACK (a test targeting a known-blank-hierarchy app drives + asserts via a self-validated CV/OCR pixel oracle, not a hierarchy-text attribute); paired §1.1 mutation makes the pixel-OCR analyzer PASS its golden-bad fixture → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md §11.4.117 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no --hierarchy-is-content-oracle, --skip-pixel-fallback, --unvalidated-ocr-OK, --hardcoded-ocr-threshold-OK flag exists.

§11.4.118 — Discovery-pressure to confirm known-issue-set completeness mandate (1.1.8-dev remediation, 2026-06-03)

Short tag: discovery-pressure-completeness.

Forensic anchor (genericised, 2026-06-03). A remediation cycle has a list of reported defects. Fixing exactly those is NECESSARY but NOT SUFFICIENT — the reported set is what the operator happened to notice, biased by what they tested (“we only see what we test”). The remediation ran parallel discovery + stress streams across ALL target devices, exercising subsystems the reported defects did NOT touch, to CONFIRM the reported set IS the complete critical set (and to surface any unreported defect before the user does). The discovery is anti-bluff only if it provides PROVABLE coverage: a list of the subsystems / journeys actually exercised, each with the outcome (no-new-issue vs new-issue-filed), so “we found nothing else” is backed by “here is everything we looked at,” not by “we stopped looking.”

The mandate. A remediation / release cycle MUST NOT treat “every reported defect is fixed” as “the build is good.” After (or in parallel with) fixing the reported set, the cycle MUST run a discovery + stress pass across ALL target devices / environments that deliberately exercises subsystems, journeys, and edge cases BEYOND the reported defects — to confirm the reported set is the COMPLETE critical set and to surface unreported defects before the end user does. The discovery pass MUST produce PROVABLE coverage: an enumerated list of the subsystems / user-journeys / stress scenarios actually exercised, each with its outcome (no-new-issue, or a new tracker entry filed per §11.4.15 / §11.4.16). “We found no other issues” is a §11.4 bluff unless accompanied by “here is the enumerated set we

exercised” — absence of evidence of looking is not evidence of absence. Newly-discovered defects trigger §11.4.4 test-interrupt + the §11.4.114/§11.4.115 isolation→RED→fix loop.

Why necessary-not-sufficient. A cycle that fixes only the reported defects ships with the operator’s coverage bias baked in: the subsystems the operator did not test on the broken build are still unverified on the fixed build. Discovery pressure (everyday-user-simulation journeys + chaos/stress per §11.4.85, run autonomously per §11.4.52 across every device per §11.4.119) is how the cycle earns the claim “this is the complete critical set,” converting an unbounded “are there other bugs?” into a bounded, evidenced “here is what we exercised and what we found.”

Honest boundary (§11.4.6). Discovery pressure reduces the unknown-unknown surface; it does not prove zero remaining defects (no finite test set can). The claim the cycle earns is “the reported set + the enumerated discovery set are all addressed,” with the discovery set’s coverage explicitly bounded — never “the build is bug-free.” Subsystems NOT exercised are stated as such (an honest coverage gap per §11.4.3), not silently implied clean.

Classification: universal (§11.4.17) — discovery-pressure to confirm known-issue-set completeness is a platform-neutral release-confidence discipline reusable by ANY project with a remediation/release cycle; the consuming project supplies its subsystem inventory, everyday-user-journey set, stress scenarios, and target-environment matrix per §11.4.35.

Composes with §11.4.4 (test-interrupt — a discovery finding STOPS the cycle), §11.4.5 / §11.4.69 (the discovery pass captures evidence per subsystem exercised), §11.4.25 (full-automation-coverage — the discovery set is part of the coverage ledger), §11.4.40 (full-suite retest — discovery runs alongside / inside it), §11.4.42 (iteration-discipline — discovery is the smoke-then-full progression’s breadth dimension), §11.4.52 (autonomous-validation — discovery journeys run without a human), §11.4.85 (stress + chaos — discovery pressure IS the stress dimension applied for completeness), §11.4.114 / §11.4.115 (a discovered defect enters the isolation→RED→fix loop), §11.4.119 (single-resource-owner — parallel per-device discovery needs the ownership partition).

Propagation. Propagation gate CM-COVENANT-114-118-PROPAGATION enforces the literal anchor 11.4.118 across the consumer fleet; paired §1.1 meta-test mutation strips the literal → the gate FAILs. Recommended per-family gate CM-DISCOVERY-COVERAGE-ENUMERATED (a remediation/release cycle records an enumerated subsystem/journey discovery-coverage list with per-item outcome, not a bare “no other issues found”); paired §1.1 mutation replaces the enumerated list with a bare no-issues claim → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md §11.4.118 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no --reported-set-suffices, --skip-discovery-pass, --no-issues-without-coverage-list, --assume-complete flag exists.

§11.4.119 — Single-resource-owner partitioning for parallel hardware testing mandate (1.1.8-dev remediation, 2026-06-03)

Short tag: single-resource-owner-partitioning.

Forensic anchor (genericised, 2026-06-03). Multiple parallel test/discovery streams ran against shared physical hardware (devices D3 + D4, an HDMI audio sink). Two streams driving media playback on the SAME device at once corrupt each other's evidence — overlapping audio routes the sink reports the wrong codec, concurrent `am_start` fights over the foreground, a second stream's input events land in the first stream's app. The remediation enforced: exactly ONE stream OWNS each device's exclusive/media resource at a time; other streams targeting that device are read-only (probes / dumpsys / log capture). Parallelism is partitioned BY DEVICE — stream A owns D3's media, stream B owns D4's media — so the parallel-development pipeline (§11.4.58 / §11.4.103) runs without the streams clobbering each other's captured evidence.

The mandate. When multiple parallel work/test/discovery streams exercise SHARED hardware or any exclusive-access resource (a media-playback path, a single HDMI/audio sink, an exclusive device handle, a single serial/JTAG line, a GPU under exclusive capture), exactly ONE stream MUST own each such resource at a time. The exclusive owner drives the resource (playback, input injection, capture); every other concurrent stream targeting the same resource MUST be READ-ONLY (passive probes — dumpsys / `/proc` / `/sys` reads / sink-side network probes / log tails) for the duration. Parallelism MUST be partitioned by resource: distinct devices / sinks / handles run fully concurrent (stream-per-device), but the same device's exclusive resource is single-owner. The ownership MUST be enforced by an advisory lock / token (per §11.4.58 L3 contention-path locking + §11.4.84 working-tree quiescence's hardware analogue), not by convention — a second stream MUST be unable to silently start media on a device another stream owns. Ownership is event-driven: a stream claims a device when its exclusive resource frees and releases it the moment its work completes, so the next queued stream can claim it (per §11.4.103 auto-backfill).

Why single-owner. Concurrent drivers of one exclusive resource produce CROSS-CONTAMINATED evidence: the sink reports whichever stream's audio won the race, the foreground belongs to whichever `am_start` landed last, input events interleave. A PASS captured under contention is a §11.4 evidence-integrity bluff — the captured artefact does not reflect what the test under inspection actually did. Single-owner partitioning is the precondition that makes per-stream captured evidence trustworthy under parallelism.

Honest boundary (§11.4.6). A read-only probe stream is genuinely read-only — it MUST NOT issue any state-changing command (no `am_start`, no input injection, no `setprop`, no playback) against a device it does not own; a “read-only” stream that mutates state is the contention the partition forbids. When the resource genuinely cannot be partitioned (a single device the whole batch needs to drive), the streams MUST serialize on it (single-owner over time), not run concurrently and hope.

Classification: universal (§11.4.17) — single-owner-per-exclusive-resource partitioning is a platform-neutral parallel-testing discipline reusable by ANY project with parallel streams contending over shared hardware / exclusive handles; the

consuming project supplies its resource inventory (which devices / sinks / handles are exclusive) and lock mechanism per §11.4.35. Strict refinement of §11.4.58 / §11.4.103 for the hardware-contention case.

Composes with §11.4.5 / §11.4.69 (captured-evidence integrity — single-owner is the precondition for trustworthy per-stream evidence), §11.4.13 (sink-side evidence — a shared sink is the canonical single-owner resource), §11.4.50 (deterministic consistency — contended evidence is non-deterministic, defeating N-iteration stability), §11.4.58 (parallel-development PWU — §11.4.119 is the hardware-resource refinement of its L3 contention-path locking), §11.4.82(F) (parallel multi-device testing — per-device streams ARE the partition), §11.4.84 (working-tree quiescence — the hardware analogue: a stream owns the device's exclusive state cleanly before it drives it), §11.4.103 (continuous parallel streams — single-owner + event-driven release enables the auto-backfill), §11.4.118 (discovery pressure — parallel per-device discovery needs this partition).

Propagation. Propagation gate CM-COVENANT-114-119-PROPAGATION enforces the literal anchor 11.4.119 across the consumer fleet; paired §1.1 meta-test mutation strips the literal → the gate FAILs. Recommended per-family gate CM-SINGLE-RESOURCE-OWNER-PARTITION (parallel hardware-test streams acquire a per-resource ownership lock before any state-changing command; concurrent same-resource streams are read-only); paired §1.1 mutation removes the ownership-lock acquisition (letting two streams drive one device) → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md §11.4.119 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no --allow-concurrent-resource-drivers, --skip-ownership-lock, --read-only-may-mutate, --contended-evidence-OK flag exists.

§11.4.120 — Fix-breaks-its-own-gate reconciliation mandate (1.1.8-dev remediation, 2026-06-03)

Short tag: gate-reconciliation-not-fake-pass.

Forensic anchor (genericised, 2026-06-03). A correct fix removed the regression-causing behaviour (the audio fix removed the handledByDualHdmi short-circuit that diverted 5.1 off the multichannel thread). But an EARLIER gate had been authored to assert the OLD — now-known-broken — behaviour's presence. After the fix, that gate FAILED. The FAIL was the CORRECT signal that the fix worked: the gate was asserting a behaviour the fix deliberately removed. Two wrong responses are forbidden — fake-passing the gate (editing it to always pass), or reverting the correct fix to satisfy the stale gate. The right response is to RECONCILE the gate: rewrite it to assert the NEW mechanism (evidence-backed), so it once again guards the now-correct behaviour.

The mandate. When a correct fix causes a pre-existing gate / test to FAIL because that gate asserted the OLD (now-removed-or-changed) behaviour, the gate FAIL is the CORRECT signal that the fix landed — it MUST NOT be suppressed by either of the two forbidden responses: (1) **fake-passing** the gate — editing it to always pass, weakening its assertion to a tautology, or deleting it to make the FAIL

disappear (a §11.4 bluff at the gate layer, and a §11.4.84 mutation-residue risk); (2) **reverting the correct fix** to satisfy the stale gate (re-introducing the defect to keep a now-wrong gate green). The required response is **RECONCILIATION**: rewrite the gate so it asserts the NEW mechanism the fix introduced, backed by captured evidence of the new correct behaviour, so the gate once again guards the right invariant. After reconciliation, the gate's paired §1.1 mutation **MUST** be updated too (the mutation must break the NEW invariant). The reconciliation **MUST** be a visible, evidence-cited change — “rewrote gate X to assert because fix Y removed ; new captured evidence: ” — never a silent assertion-weakening.

Distinguish reconciliation from bluffing. Reconciliation rewrites the gate to assert a REAL new invariant (the fix's intended mechanism), proven by captured evidence, with the paired mutation updated to break it. Bluffing weakens the gate so it can no longer FAIL on the defect class it was meant to catch. The discriminator: after reconciliation the gate + its mutation still form a valid §1.1 pair (mutate the new invariant → gate FAILs); a bluffed gate's mutation no longer makes it FAIL (the assertion became a tautology). A reconciled gate is provably still a guard; a bluffed gate is provably blind.

Honest boundary (§11.4.6). A gate FAIL after a fix is NOT automatically “the gate is stale, reconcile it.” It **MUST** first be investigated per §11.4.102 — the FAIL may be the gate correctly catching a REGRESSION the fix introduced (the fix broke something else), in which case the FIX is wrong, not the gate. Reconcile **ONLY** when investigation **PROVES** the gate asserted old-correct-now-removed behaviour **AND** the new behaviour is the intended, evidence-confirmed mechanism. “The gate is just stale” without that investigation is a §11.4.6 guess that can mask a real regression.

Classification: universal (§11.4.17) — reconcile-the-gate-don't-fake-pass-and-don't-revert is a platform-neutral fix-vs-gate-conflict discipline reusable by ANY project with regression gates; the consuming project supplies its gate framework and paired-mutation mechanism per §11.4.35.

Composes with §11.4.1 (FAIL-bluffs — fake-passing the gate is the inverse bluff), §11.4.4 (test-interrupt — a post-fix gate FAIL triggers investigation before reconcile), §11.4.6 (no-guessing — “stale gate” is a hypothesis until investigation proves it), §11.4.84 (working-tree quiescence — a gate edited to always pass is mutation residue that must never ship), §11.4.102 (systematic-debugging — the mandatory investigation that distinguishes stale-gate from new-regression), §11.4.108 (four-layer fix-verification — the reconciled gate asserts the fix's runtime signature / new mechanism), §1.1 (paired mutation — the reconciled gate's mutation must break the NEW invariant, re-proving it is a guard not a tautology).

Propagation. Propagation gate CM-COVENANT-114-120-PROPAGATION enforces the literal anchor 11.4.120 across the consumer fleet; paired §1.1 meta-test mutation strips the literal → the gate FAILs. Recommended per-family gate CM-GATE-RECONCILED-NOT-FAKE-PASSED (when a fix changes a behaviour a gate asserted, the gate's reconciliation is an evidence-cited rewrite whose paired mutation still makes it FAIL — never an assertion weakened to a tautology); paired §1.1 mutation weakens a reconciled gate to always pass and asserts its mutation no longer makes it FAIL → the meta-gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md §11.4.120 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--fake-pass-stale-gate`, `--revert-fix-for-gate`, `--weaken-assertion-to-pass`, `--delete-failing-gate` flag exists.

§11.4.121 — No-commit-while-build-writes-tracked-artifacts mandate (1.1.8-dev remediation, 2026-06-03)

Short tag: `no-commit-during-artifact-write`.

Forensic anchor (genericised, 2026-06-03). The build pipeline’s per-component build steps write produced artifacts (forked-app APKs) INTO tracked directories that are under version control (`prebuilt_apps/`). Running `git add -A` / a commit while those steps are mid-write stages a PARTIAL or stale artifact — a race between the build writing the file and git reading it — so the commit captures a half-written or pre-rebuild artifact. The remediation deferred the commit to BUILD-COMPLETION: the tree is quiescent (no build step writing tracked files) AND the freshly-built artifacts are captured whole (no stale prebuilt in git). This is the artifact-write analogue of §11.4.84 (no commit while a mutation gate is in flight) at the build-output layer.

The mandate. A commit (especially `git add -A` / any broad stage) MUST NOT run while a build / packaging / generation step is actively writing artifacts INTO tracked directories. Doing so races the writer and stages a partial or stale artifact — the commit captures bytes that are mid-write or pre-rebuild, a §11.4.108 SOURCE→ARTIFACT integrity failure landed in version control. The commit MUST be deferred until the build step that writes tracked artifacts has COMPLETED, so that: (1) the working tree is quiescent at the artifact layer (no step writing a tracked file); (2) the committed artifacts are the FRESH, whole outputs of the completed build (no stale pre-rebuild artifact, no half-written file). Before committing tracked build outputs, the committer MUST verify the writing step finished (process exit, completion marker, or per-artifact `mtime ≥ build-start`) — a build still in flight that writes tracked dirs is a HOLD on the commit, not a race to win.

Composition with §11.4.84. §11.4.84 forbids committing while a mutation gate is in flight in the checkout (mutation residue). §11.4.121 is the BUILD-OUTPUT analogue: forbid committing while a build is in flight WRITING TRACKED ARTIFACTS (partial/stale artifact residue). Both close the same class of defect — a commit captures transient, non-final tree state — at two different write-sources (mutation experiments vs build outputs). Where build outputs land OUTSIDE version control (`gitignored out/ / dist/`), the race does not apply; the mandate binds specifically to build steps writing into TRACKED directories.

Honest boundary (§11.4.6). “The build probably finished” is not “the build finished” — verify with a completion signal (exit code / marker / mtime), not an assumption, before committing tracked artifacts. Committing source changes that DON’T touch the build’s tracked-artifact directories is fine mid-build (those files are not being written by the build); the hold applies only to the tracked dirs the in-flight build writes.

Classification: universal (§11.4.17) — deferring commit until a tracked-artifact-writing build completes is a platform-neutral commit-integrity discipline reusable by ANY project whose build writes produced artifacts into version-controlled

directories; the consuming project supplies which build steps write which tracked directories and its completion-signal mechanism per §11.4.35. The PROJECT-SPECIFIC instance (which tracked directory — e.g. a prebuilt-apps directory — and which build steps write it) is recorded in the consuming project’s own governance per §11.4.35.

Composes with §11.4.6 (no-guessing — verify build completion, never assume), §11.4.30 (no-versioned-build-artifacts — artifacts that ARE legitimately tracked, like signed prebuilts, are the exact case this mandate protects from partial commit; gitignored build outputs are exempt), §11.4.58 / §11.4.103 (parallel-development — the commit/merge stage is serialized AFTER the build stage for this reason), §11.4.84 (working-tree quiescence — §11.4.121 is its build-output analogue), §11.4.88 (background-push — commit timing is decoupled from push, but the commit itself still waits for artifact-write completion), §11.4.108 (four-layer fix-verification — a partial/stale committed artifact is the ARTIFACT-layer integrity failure this prevents from entering version control), §11.4.96 (safe-parallel-work-with-long-build — committing tracked build outputs mid-build is an UNSAFE operation in its catalogue).

Propagation. Propagation gate CM-COVENANT-114-121-PROPAGATION enforces the literal anchor 11.4.121 across the consumer fleet; paired §1.1 meta-test mutation strips the literal → the gate FAILs. Recommended per-family gate CM-NO-COMMIT-DURING-ARTIFACT-WRITE (the commit wrapper refuses to stage tracked build-artifact directories while a build step writing them is in flight — checks a build-in-progress lock / completion marker); paired §1.1 mutation removes the build-in-flight check from the commit wrapper → the gate FAILs. (Gate-code implementation lands as a separate work item; this anchor defines the contract.)

Canonical authority: this Constitution.md §11.4.121 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no --commit-during-build, --stage-partial-artifact-OK, --assume-build-finished, --skip-build-completion-check flag exists.

§11.4.122 — No-silent-removal-of-existing-components-without-operator-confirmation mandate (User mandate, 2026-06-03)

Forensic anchor — verbatim user mandate (2026-06-03):

“Never ever remove any application, system component or service from already existing codebase / System without interactively asked question to us! THIS IS MANDATORY RULE / CONSTRAINT!”

Forensic case study (this project, FACT — captured). During the 1.1.8-dev burn-down, two shipped capabilities — F2 (an Apple-TV-class application) and F4 (a Huawei HMS / Mobile-Services component) — were removed from the existing System WITHOUT first asking the operator. The operator reversed both removals. A removal that the operator has to discover and reverse after the fact is a defect of the same severity class as a §11.4 PASS-bluff: the System silently lost a user-facing capability the operator never agreed to drop.

The mandate. No application, system component, service, package, feature, driver, module, library, prebuilt asset, or any other already-existing capability of the existing codebase / shipped System may be removed (deleted, dropped from the package set, disabled-into-non-shipping, un-bundled, de-listed, or otherwise made unavailable to the end user) WITHOUT FIRST interactively asking the operator and receiving an EXPLICIT keep-or-remove decision. The question MUST be posed through the platform's interactive clarification mechanism per §11.4.66 (AskUserQuestion on Claude Code; the documented interactive prompt on other runtimes) — NEVER a free-text “should I remove X?” buried in narrative, NEVER a silent removal justified post-hoc, NEVER an autonomous removal decision. A silent removal is a **release blocker** regardless of how well-intentioned the rationale (deduplication, “it was broken anyway”, “geo-restricted”, “incompatible”, “superseded”) — the operator decides, the agent asks.

What counts as a removal (non-exhaustive). Deleting an app/APK/binary from the build's package set (e.g. a project's PRODUCT_PACKAGES / device.mk / equivalent manifest), removing a service from the init/boot/service-registry set, dropping a kernel module / driver / config from the shipping configuration, un-bundling a prebuilt asset, deleting a submodule or a submodule's shipped output, removing a feature flag that gated a live capability, or any edit whose NET EFFECT is “an end-user capability that shipped before no longer ships.” Adding, replacing-with-equivalent-or-better (where the operator has approved the replacement), or fixing a capability is NOT a removal. When uncertain whether an edit constitutes a removal, treat it AS a removal and ask (per §11.4.6 no-guessing + §11.4.101 block-only-when-truly-needed — removal of an existing user-facing capability is high-blast-radius and MUST be operator-confirmed, never autonomously decided).

Composition. §11.4.122 composes with §11.4.66 (interactive-clarification — the removal question is asked via the platform interactive mechanism, not free-text), §11.4.101 (autonomous-decision-over-blocking — a removal of an existing user-facing capability is exactly the irreversible-ish, high-blast-radius class §11.4.101 RESERVES for an operator question rather than an autonomous decision; §11.4.122 makes the “ask the operator” outcome mandatory for this class, NOT optional), §11.4.90 (Obsolete classification — the tracked path to actually DROP a capability is: ask the operator → operator approves → mark the item Obsolete (→ Fixed.md) with Obsolete-Details reason feature-removed AND an operator-approval citation → then remove; the removal never precedes the operator's yes), §11.4.112 (structural-impossibility won't-fix — even a capability proven structurally impossible to keep is closed via the §11.4.90 path with operator awareness, not silently deleted), §11.4.6 (no-guessing — “the operator probably wants this gone” is the exact guess forbidden), §11.4.40 / §11.4.42 (full-suite-retest / iteration-discipline — a removal landed without the operator's yes fails the release gate).

Enforcement. Propagation gate CM-COVENANT-114-122-PROPAGATION enforces the literal anchor 11.4.122 across the consumer fleet (every CLAUDE.md / AGENTS.md / QWEN.md). Recommended gate CM-NO-SILENT-COMPONENT-REMOVAL (a diff-driven detector that flags a removal of a shipped capability — a deletion from the project's package/service/module manifest set — without a matching operator-approval citation in the commit / tracker entry, per the §11.4.110 change-impact-clash pattern; a removal whose tracker item is Obsolete (→ Fixed.md) with reason feature-removed + operator-approval citation passes). Paired §1.1 meta-test mutations: strip the 11.4.122 literal → propagation gate FAILs; introduce a manifest removal with no operator-approval citation → CM-NO-SILENT-COMPONENT-REMOVAL FAILs. (Gate-code implementation lands as a separate follow-up work item; this anchor defines the contract.)

Classification: universal (§11.4.17) — “never silently remove an existing end-user capability without operator confirmation” is a platform-neutral discipline reusable by ANY project that ships a set of user-facing capabilities; the consuming project supplies its concrete capability-manifest paths (package set, service registry, module config) per §11.4.35.

Canonical authority: this Constitution.md §11.4.122 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--remove-without-asking`, `--silent-removal`, `--autonomous-removal-OK`, `--dedup-removal-exempt`, `--it-was-broken-anyway` flag exists.

§11.4.123 — Rock-solid-proof-or-deep-research mandate (User mandate, 2026-06-03)

Forensic anchor — verbatim user mandate (2026-06-03):

“Every single reported issue MUST BE fully and 100% validated with rock solid proofs! Nothing can be considered fixed or completed without hard evidence! No false results or bluff(s) of any kind is allowed! If we are not sure on how to achieve full testing, validation and verification of something we MUST ALWAYS perform deep web research for all possible data (articles, documentation, guides, and other resources) and opensourced codebases which we can use to solve our problems and perform testing with validation and verification which produces rock-solid evidence(s) and leaves no space for false results or any kind of bluff!”

Forensic case study (FACT — captured). In the 1.1.8-dev remediation cycle the validation method for two feature classes was, at first, genuinely unclear: relocating a FLAG_SECURE secure surface to a secondary display (whose pixel capture returns black) and asserting on-screen content in non-introspectable streaming-app UIs (whose accessibility hierarchy is blank). Rather than declaring those classes “untestable” or accepting a metadata-only PASS, the cycle performed deep web research (docs/research/testing_frameworks_20260603/) into platform/HDCP behaviour, computer-vision template-matching, ROI OCR with confidence floors, freeze-detection / frame-advance liveness oracles, and network-side sink probes. That research yielded the CV/OCR/liveness/sink-probe oracle stack (now §11.4.107 + §11.4.112 + §11.4.117) that made rock-solid, evidence-producing validation possible where it had appeared impossible. The lesson — “unclear how to validate” is a research trigger, NEVER a bluff licence — is now constitutional.

The mandate. Every single reported issue, every fix, and every claimed completion MUST be fully and 100% validated with rock-solid CAPTURED proof per §11.4.5 / §11.4.69 / §11.4.107 before it may be marked fixed / implemented / completed (the §11.4.33 terminal closure vocabulary). Hard, captured evidence is the bar — nothing may be considered fixed or complete without it. Metadata-only PASS, configuration-only PASS, absence-of-error PASS, and grep-without-runtime PASS are all forbidden (§11.4 / §11.4.1) — no false results and no bluff of any kind, at any layer, are allowed.

The research-or-don’t-bluff rule (the operative addition). When the agent is UNSURE how to fully test, validate, or verify something — when no obvious evidence-producing method exists, OR the candidate method would yield only

metadata/config/absence-of-error evidence (and is therefore a bluff per §11.4.1) — the agent MUST ALWAYS first perform deep web research per §11.4.8 + §11.4.99 (official documentation, technical articles, guides, vendor references, standards, issue trackers, and open-source codebases the project can reuse or adapt) to DISCOVER or BUILD a validation method that produces rock-solid evidence and leaves no space for a false result. Declaring something “untestable”, “not automatable”, or accepting a metadata-only PASS WITHOUT having first exhausted this deep-research path is itself a §11.4.123 violation — the same severity class as a PASS-bluff. The research output (cited source URLs + the evidence-producing method it yielded, OR the literal “NO external solution found — original work” per §11.4.8) is the captured proof that the path was exhausted. Only after that research genuinely fails to yield any evidence-producing method may the item be classified PENDING_FORENSICS: / Operator-blocked (§11.4.21) / structurally-impossible won’t-fix (§11.4.112) — with the cited research as the evidence that the classification is earned, never a convenience.

Composition. §11.4.123 composes with §11.4.5 (captured-evidence quality is the proof shape), §11.4.6 (no-guessing — “probably fixed” without captured evidence is forbidden; the closure must be FACT), §11.4.8 (deep-web-research-before-implementation — §11.4.123 extends it from “before designing a fix” to “whenever the VALIDATION method is unclear”), §11.4.52 (autonomous-validation — the research must yield an autonomous, evidence-producing path wherever one is discoverable), §11.4.69 (universal sink-side positive-evidence taxonomy — every PASS cites a taxonomy-class artefact), §11.4.99 (latest-source verification of the discovered method), §11.4.107 (the AV/test-validation oracle techniques the 1.1.8-dev research yielded), §11.4.118 (discovery-pressure — the same “absence of evidence of looking is not evidence of absence” principle applied to per-issue validation). It also composes with §11.4.21 / §11.4.112 (the earned escape valves — reachable ONLY after the deep-research path is exhausted with cited evidence).

Enforcement. Propagation gate CM-COVENANT-114-123-PROPAGATION enforces the literal anchor 11.4.123 across the consumer fleet (every CLAUDE.md / AGENTS.md / QWEN.md). Recommended gate CM-ROCK-SOLID-PROOF-OR-RESEARCH (a closure-audit detector that flags any item closed Fixed/Implemented/Completed whose tracker entry lacks a captured-evidence artefact path per §11.4.69, AND any PENDING_FORENSICS:/Operator-blocked/structurally-impossible classification whose entry lacks a cited deep-research trail per §11.4.8/§11.4.99). Paired §1.1 meta-test mutations: strip the 11.4.123 literal → propagation gate FAILs; close an item with a metadata-only PASS and no captured-evidence path → CM-ROCK-SOLID-PROOF-OR-RESEARCH FAILs; classify an item Operator-blocked with no cited research trail → the same gate FAILs. (Gate-code implementation lands as a separate follow-up work item; this anchor defines the contract.)

Classification: universal (§11.4.17) — “every closure carries rock-solid captured proof, and unclear-how-to-validate triggers deep research rather than a bluff” is a platform-neutral discipline reusable by ANY project; the consuming project supplies its concrete capture mechanisms + research corpora per §11.4.35.

Canonical authority: this Constitution.md §11.4.123 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no --metadata-pass-suffices, --skip-proof, --untestable-without-research, --config-only-closure-OK, --bluff-when-unsure flag exists.

§11.4.124 — Dead/unwired-code investigate-before-remove mandate (User mandate, 2026-06-04)

Forensic anchor — verbatim user mandate (2026-06-04):

“Before removing any seemingly-dead (zero-importer / unwired) codebase, we MUST investigate via git history where/how it was originally used and how it became dead. Removal is permitted ONLY when we have captured PROOF it is genuinely no longer needed — and that removal MUST be its own separate commit with a proper descriptive message. If there is no such proof, the code MUST be investigated for where/how it should be wired in properly, and any missing or unwired tests MUST be added. We MUST ALWAYS be extra careful with any codebase removal.”

The defect class. “This has zero importers / is never called / is unwired, so it is dead — delete it” is a guess (§11.4.6), not a finding. A `git grep / static-analysis “no references”` result proves only *current* non-reference; it does NOT prove the code is genuinely unneeded. Code becomes apparently-dead for several distinct reasons — its call-site was deleted by mistake (a regression: the code SHOULD still be wired in), its wiring was never finished (it is intended but unwired — the fix is to WIRE it, plus its tests), it is reflectively / dynamically / build-tag / plugin / DI-container referenced (the static tool simply cannot see the reference), or it is genuinely obsolete and safe to remove. Deleting it without first distinguishing which case applies risks silently removing a needed-but-temporarily-unwired capability — the same severity class as the §11.4.122 silent-removal defect and a §11.4 PASS-bluff at the maintenance layer.

The mandate. Before removing ANY seemingly-dead codebase element (a zero-importer / never-called / unwired / unreferenced function, method, type, file, module, package, asset, config, or build target), the agent MUST FIRST investigate via git history (`git log --follow`, `git log -S<symbol> / -G<regex> pickaxe` across all history, `git log --diff-filter`, blame on the deletion of the former call-site) to establish, as captured FACT (§11.4.6): (1) WHERE and HOW it was originally used / wired in; (2) WHEN and HOW it became dead — its call-site removed deliberately, removed by mistake (regression), never completed, or made unreachable by a refactor; (3) whether the “no references” result is real OR an artefact of a reference the static tool cannot see (reflection / dynamic dispatch / build-tags / codegen / DI / plugin registry / cross-language FFI / config-driven wiring). The investigation output (the cited commits + the determination) is the captured evidence.

Removal is conditional — proof-of-genuinely-unneeded REQUIRED. Removal is permitted ONLY when the investigation produces captured PROOF the element is genuinely no longer needed (its purpose is fully superseded, its former caller is intentionally and permanently gone, no hidden reference exists). When that proof exists, the removal MUST be its OWN SEPARATE COMMIT (not bundled with unrelated work — so it is independently reviewable and revertible, composing with §11.4.84 working-tree quiescence + §11.4.92 multi-pass) carrying a proper descriptive message that cites the git-history investigation evidence (the commits that wired it in, the commit that orphaned it, the determination that it is unneeded) and — where the element is an end-user-facing capability — the §11.4.122 operator-confirmation. The §11.4.90 tracked path applies: classify the tracker item `Obsolete` (→ `Fixed.md`) with the appropriate reason.

No proof ⇒ **wire it in + add the missing tests (do NOT delete)**. If the investigation does NOT yield proof the element is unneeded, the element MUST NOT be removed. Instead it MUST be investigated for WHERE and HOW it should be wired in properly (restore the deleted call-site if its removal was a regression per §11.4.114; complete the intended wiring if it was never finished), AND any missing or unwired tests for it MUST be added (the absence of a test is itself part of why it drifted into apparent-deadness — composes §11.4.27 100%-test-type-coverage + §11.4.43 TDD-fix + §11.4.115 RED-baseline). Apparent-deadness with no removal-proof is a wiring/coverage GAP to close, never a delete-on-sight.

Extra-caution default. The agent MUST ALWAYS be extra careful with any codebase removal: when uncertain whether removal-proof is sufficient, default to NOT removing (investigate + wire + test), per §11.4.6 no-guessing + §11.4.101 (an irreversible-ish, possibly-needed removal whose safety cannot be determined from evidence is exactly the class that blocks to the operator rather than being autonomously deleted) + §11.4.122 (removal of a live end-user capability requires operator confirmation). “Probably dead” is never sufficient; the bar is captured proof.

Composition. §11.4.124 composes with §11.4.6 (no-guessing — “zero importers therefore dead” is the exact guess forbidden; the determination must be FACT from git history), §11.4.8 (deep-research — git history IS the primary research corpus for how the code was used + how it died; external research applies when the wiring mechanism is unclear), §11.4.84 / §11.4.92 (the separate-removal-commit discipline — quiescent tree, multi-pass evaluation, independently revertible), §11.4.90 (Obsolete classification is the tracked path for a proven removal), §11.4.101 (autonomous-decision-over-blocking — a removal whose safety cannot be proven blocks rather than proceeds), §11.4.114 (last-known-good-tag isolation — a call-site deleted by mistake is a regression to repair, the same diagnostic the “became dead by mistake” case demands), §11.4.122 (no-silent-removal — §11.4.124 is its complement: §11.4.122 governs removal of KNOWN-shipped capabilities by operator confirmation, §11.4.124 governs removal of SEEMINGLY-dead code by git-history proof; an end-user-facing seemingly-dead element triggers BOTH), §11.4.27 / §11.4.43 / §11.4.115 (the missing-tests obligation when the element is wired in rather than removed).

Enforcement. Propagation gate CM-COVENANT-114-124-PROPAGATION enforces the literal anchor 11.4.124 across the consumer fleet (every CLAUDE.md / AGENTS.md / QWEN.md). Recommended gate CM-DEAD-CODE-INVESTIGATE-BEFORE-REMOVE (a commit-audit detector: any commit whose diff is a NET DELETION of a code element — function/type/file/module/package/asset/build-target — MUST be a removal-only commit whose message cites the git-history investigation evidence, OR the deletion is part of a tracked Obsolete (→ Fixed.md) item per §11.4.90; a net-deletion commit bundled with unrelated changes OR lacking the investigation citation FAILs). Paired §1.1 meta-test mutations: strip the 11.4.124 literal → propagation gate FAILs; introduce a net-deletion commit with no investigation citation and not part of an Obsolete item → CM-DEAD-CODE-INVESTIGATE-BEFORE-REMOVE FAILs. (Gate-code implementation lands as a separate follow-up work item; this anchor defines the contract.)

Classification: universal (§11.4.17) — “investigate git history + capture proof-of-unneeded before removing seemingly-dead code, otherwise wire it in and add its tests; removal is its own descriptive commit” is a platform-neutral maintenance

discipline reusable by ANY codebase; the consuming project supplies its concrete static-analysis / importer-graph tooling and its hidden-reference mechanisms (reflection, DI, build-tags, codegen) per §11.4.35.

Canonical authority: this Constitution.md §11.4.124 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--zero-importers-means-dead`, `--delete-unwired-on-sight`, `--skip-git-history-investigation`, `--remove-without-proof`, `--bundle-removal-with-other-work` flag exists.

§11.4.125 — Code-review-agent gate before pre-build + main build (mandatory multi-layer review) (User mandate, 2026-06-04)

Forensic anchor — verbatim user mandate (2026-06-04):

“After all fixes/changes/implementations are done, BEFORE running pre-build tests and the main build, dispatch code-review agent(s) that analyze all work done + all existing data/facts + the existing codebase + current git history to determine quality, safety, and whether the fixes/changes will REALLY work; they MUST validate and verify that every test covering the fixes/changes genuinely validates the work with NO chance of false results or bluff of any kind. Any finding MUST be fixed, polished, improved, and covered with additional tests before the build proceeds. Multiple strong layers of checks.”

The defect class. Running the pre-build gate sweep and the main (artifact) build directly after a batch of fixes/changes — without an independent review pass first — lets a defective fix, an error-prone change, a bluff test, or a test that does not genuinely validate the work-under-test consume an expensive build cycle and (worse) ship green. Pre-build gates and the build prove the codebase is internally consistent + buildable; they do NOT prove the work is correct, safe, will-really-work for the end user, or that the tests covering it actually catch its negation. That assurance is an independent-review responsibility that MUST run BEFORE the build, not after.

The mandate. After all fixes / changes / implementations in a batch are done, and BEFORE running the pre-build test sweep and the main build (for ANY project), the agent MUST dispatch one or more dedicated code-review agent(s) (subagent-driven by default per §11.4.70 / §11.4.20) that perform a multi-layer review. The review MUST: (1) **analyze ALL work done in the batch** — every fix / change / implementation, its source diff, and its stated intent; (2) **analyze ALL existing data + facts** — captured evidence per §11.4.5 / §11.4.69 / §11.4.107, tracker entries, prior findings, the §11.4.108 runtime-signature registry; (3) **analyze the existing codebase** — the blast radius of each change (§11.4.92 multi-pass), cross-feature interaction, contract integrity of every dependency; (4) **analyze current git history** — what each change touched, how it composes with concurrent / recent work, whether it reproduces a known-broken pattern (§11.4.114 / §11.4.124 git-history discipline); (5) **determine quality, safety, and will-it-REALLY-work** — assess whether each change is robust + not error-prone (will not solve A while creating B), safe (no host-safety / data-safety / security regression), and genuinely delivers the end-user-visible behaviour it claims (§11.4 / §107); (6) **validate + verify the tests covering the work** — confirm every test that covers a fix/change genuinely

exercises the work-under-test and catches its negation, with ZERO chance of a false result or bluff of any kind (a test that PASSES on broken-for-the-user work, a metadata-only / config-only / absence-of-error / grep-without-runtime assertion, or a gate whose paired §1.1 mutation does not make it FAIL is a finding the review MUST surface per §11.4.1 / §11.4.6 / §11.4.107 / §1.1).

Findings block the build. Any finding the review surfaces — a defect, an error-prone change, a safety risk, a will-not-really-work conclusion, a bluff/false-result-capable test, a missing-coverage gap — MUST be fixed, polished, improved, and covered with additional tests (four-layer per §11.4.4(b), TDD-RED-first per §11.4.43 / §11.4.115) BEFORE the pre-build sweep and main build proceed. The review MUST iterate (re-review after each remediation) until it surfaces no remaining blocking findings; only then does the build proceed. The review is itself anti-bluff: its conclusions are captured evidence (§11.4.5 / §11.4.69), and a review that rubber-stamps a defective batch is the same severity class as a PASS-bluff.

Multiple strong layers. The review is one of multiple strong layers of checks, complementing — never replacing — the pre-build gate sweep (§1 four-layer), the §11.4.92 multi-pass change-evaluation, the §11.4.108 four-layer fix-verification, the §11.4.110 pre-build build-readiness verdict + clash detector, and the post-build / runtime-on-clean-target / user-visible layers. §11.4.125 sits at the seam between “all batch work done” and “start the pre-build + build”: it is the independent-review gate that the other layers do not provide (the multi-pass per §11.4.92 is self-review by the change author; §11.4.125 adds the structurally-separated reviewer seam per §11.4.70).

Composition. §11.4.125 composes with §11.4 (the end-user-quality covenant the review enforces), §11.4.1 (FAIL-bluffs equally caught), §11.4.4 (four-layer coverage of every remediation the review demands), §11.4.6 (no-guessing — review conclusions are FACT from captured evidence, never “looks fine”), §11.4.40 (the full-suite retest is a later gate; §11.4.125 is the pre-build review gate that precedes it), §11.4.43 (TDD-fix — remediation lands RED-first), §11.4.50 (deterministic-consistency of the tests the review validates), §11.4.70 / §11.4.20 (the review IS a subagent dispatch — structurally-separated reviewer eliminates self-review blind spots), §11.4.92 (multi-pass change-evaluation is the author-side pass; §11.4.125 is the independent-reviewer pass), §11.4.102 (any defect the review finds triggers systematic-debugging before remediation), §11.4.107 (the review verifies AV/test-validation analyzers are not themselves bluffs), §11.4.108 (the review checks each fix’s runtime signature is declared + verifiable), §11.4.110 (the build-readiness verdict + clash detector are mechanical gates the review complements with judgement).

Enforcement. Propagation gate CM-COVENANT-114-125-PROPAGATION enforces the literal anchor 11.4.125 across the consumer fleet (every CLAUDE.md / AGENTS.md / QWEN.md). Recommended gate CM-CODE-REVIEW-GATE-BEFORE-BUILD (the build orchestrator records a code-review-completed marker — with the review’s findings + their remediation-or-no-blocking-findings determination — produced AFTER the batch’s last fix and BEFORE the pre-build sweep + main build; a build that starts without a fresh review marker for the current batch FAILs). Paired §1.1 meta-test mutations: strip the 11.4.125 literal → propagation gate FAILs; start a build with no fresh code-review marker for the batch → CM-CODE-REVIEW-GATE-BEFORE-BUILD FAILs. (Gate-code implementation lands as a separate follow-up work item; this anchor defines the contract.)

Classification: universal (§11.4.17) — “an independent multi-layer code-review-agent gate analyzing all work + data + codebase + git history for quality/safety/will-it-really-work + genuine test-validation, with findings blocking the build, runs BEFORE the pre-build sweep + main build” is a platform-neutral discipline reusable by ANY project that builds an artifact from source; the consuming project supplies its concrete review-marker mechanism + build-orchestrator integration per §11.4.35.

Canonical authority: this Constitution.md §11.4.125 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--skip-code-review`, `--build-without-review`, `--no-review-gate`, `--review-optional`, `--trust-the-author` flag exists.

§11.4.126 — Default autonomous-loop working mode from first prompt (User mandate, 2026-06-04)

Forensic anchor — verbatim user mandate (2026-06-04):

“Make sure that you continue work in endless fully autonomous loop, do not stop until new fully validated and verified version (tag) is created and published (all submodules and main repo) or IN A CASE OF some other main stream work until it is fully completed with all side work streams and nothing else is left in our working queue! THIS MUST BE ALWAYS the default working mode without us asking you! We tend to achieve ABSOLUTE EFFICIENCY, with this and all other projects which will incorporate this MANDATORY RULE / CONSTRAINT!!! This way of (your) working will be ALWAYS applied / followed / executed / fully respected, as soon as we assign / send first request (prompt) in the session! This stops only if we explicitly say so or nothing is left to be done in current working scope (release that will come / upcoming version)!!! Any mimicking (imitation) of this behavior / rules / mandatory constraints, false results or any kind of bluff(s) is ABSOLUTELY FORBIDDEN!!!”

The endless fully-autonomous loop is the **DEFAULT working mode**, engaged automatically the moment the operator sends the FIRST request / prompt of a session — the operator MUST NOT have to ask for it, request it, restate it, or re-enable it per session. §11.4.87 framed the endless-loop covenant as an explicit-instruction opt-in (“continue in endless loop fully autonomously” or a semantically-equivalent phrasing); §11.4.126 is the **capstone** that promotes the same covenant to always-on: from the first prompt onward, every agent operates in the §11.4.87 loop discipline as the standing default, with §11.4.94 zero-idle, §11.4.97 maximum-idle-use, §11.4.101 autonomous-decision-over-blocking, and §11.4.103 continuous-parallel-stream all engaged by default — no per-session activation handshake.

The continuation contract (the loop continues until ONE of the two terminal conditions holds): (A) **Release scope** — a new, fully-validated-and-verified version (tag) is created AND published across all owned submodules AND the main repo to all configured remotes (per §2.1 multi-upstream push + §11.4.40 full-suite-retest-before-tag + §11.4.113 absolute-no-force-push merge-onto-latest-main). (B) **Non-release main-stream scope** — the main-stream work goal is fully completed AND every side work stream is done AND the working queue holds nothing left to do for

the current scope. Until ONE of (A) or (B) holds, the agent MUST keep working — it claims the next priority item, dispatches the next parallel stream, or progresses every non-blocked item, per §11.4.42 / §11.4.72 / §11.4.94 / §11.4.103.

The loop STOPS ONLY on: (1) the operator explicitly saying so (STOP / pause / end); (2) nothing left to do in the current working scope — the upcoming release / the current main-stream goal — with the queue genuinely empty per the (A)/(B) terminal conditions; (3) a §12 host-session-safety demand (the loop yields to host safety unconditionally). Idle-while-blocked parks one work unit, it does not stop the loop — the agent keeps progressing every non-blocked item in parallel per §11.4.101 + §11.4.94 + §11.4.97.

Goal — ABSOLUTE EFFICIENCY. The discipline exists to drive the project to completion with no operator-side restart overhead, no idle gaps, no stop-and-wait round-trips. It applies to this project AND every project that incorporates this Constitution (universal per §11.4.17).

Anti-bluff — actually do the work, never imitate it. Mimicking / imitating this loop behaviour, narrating continuation without performing it, fabricating progress, or emitting false / bluff results of ANY kind is ABSOLUTELY FORBIDDEN — this composes the entire §11.4 anti-bluff covenant family (§11.4 / §11.4.1 / §11.4.2 / §11.4.5 / §11.4.6 / §11.4.50 / §11.4.69 / §11.4.107). The agent MUST genuinely perform the continuous work and capture positive evidence for every closure; a report that claims the loop ran while no real work / no captured evidence was produced is a §11.4 PASS-bluff at the operating-mode layer, the severity class this mandate specifically prohibits.

Classification: universal (§11.4.17) — a platform-neutral operating-mode discipline reusable by ANY project; the consuming project supplies its concrete release-tag / publish mechanism + working-queue source per §11.4.35. Composes with §11.4.87 (the endless-loop covenant — §11.4.126 promotes it from opt-in to always-on default) / §11.4.94 (zero-idle priority-first parallel-by-default) / §11.4.97 (maximum-idle-use + progress-update cadence) / §11.4.101 (autonomous-decision-over-blocking) / §11.4.103 (continuous parallel-stream routine) / §11.4.66 (interactive clarification ONLY when a decision is genuinely operator-blocked per §11.4.101) / §11.4.6 (no-guessing — terminal conditions determined from captured queue/tracker state, never assumed) / §11.4.40 (release-tag full-suite retest is the (A) terminal gate) / §11.4.42 (iteration discipline is the loop body) / §11.4.72 (audio-first priority within the loop) / §11.4.113 (publish via merge-onto-latest-main, no force-push) / §2.1 (multi-upstream publish) / §12 (host-session safety overrides the loop). Propagation gate CM-COVENANT-114-126-PROPAGATION (literal 11.4.126 across the consumer fleet) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: this Constitution.md §11.4.126 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no --ask-before-continuing, --single-turn-only, --not-default-loop, --mimic-OK flag exists.

§11.4.127 — Session-handoff resumption-prompt mandate (User mandate, 2026-06-06)

Forensic anchor — verbatim user mandate (2026-06-06):

“make sure that in situations like this now when new session is needed you ALWAYS prepare such sentence - which will be valid for particular moment and the phase of the project and enough for work to continue.”

When the agent determines a fresh session is needed (context-window limits, performance degradation) OR the operator asks whether a new session is needed / requests a handoff, the agent **MUST ALWAYS** prepare and provide a ready-to-paste **resumption prompt valid for that EXACT moment and project phase** — self-contained enough that pasting it into a fresh session resumes work with ZERO loss. The agent **MUST** proactively offer the resumption prompt the moment it recognises the need (not only when asked). Two variants on demand: a **SHORT** first-sentence variant (“Read <handoff docs>, then continue <terminal goal> ...”) AND a **FULL** detailed block. The prompt **MUST**: (1) point to the live handoff doc(s) — .remember/remember.md if present + docs/CONTINUATION.md per §12.10 — and instruct reading them **FIRST** + `git fetch --all`; (2) state the current **PHASE** + the immediate **NEXT** action + the terminal goal; (3) embed the exact current-state anchors (build IDs / artifact MD5, device/target serials, commit HEAD, in-flight PIDs + log paths, captured-evidence paths); (4) restate the binding constraints (anti-bluff §11.4 covenant, no-force-push §11.4.113, exact version/naming, hardware/target gotchas); (5) be **MOMENT-VALID** — reflect the actual live state, **NEVER** a generic template. The handoff doc(s) **MUST** be brought current **BEFORE** the prompt is given (§12.10). A handoff that omits the resumption prompt, or gives a stale / generic one, is a §11.4.127 violation.

Classification: universal (§11.4.17) — a platform-neutral session-continuity discipline reusable by ANY project worked by context-bounded agents; the consuming project supplies its concrete handoff-doc paths + state anchors per §11.4.35. Composes with §12.10 (CONTINUATION doc — the resumption prompt is its operator-facing entry point) / §11.4.6 (no-guessing — the embedded state is captured fact, not assumed) / §11.4.66 (interactive clarification — “do you need a new session?” is itself a §11.4.66 decision point) / §11.4.87 / §11.4.103 / §11.4.126 (the fresh session **CONTINUES** the always-on autonomous loop — the resumption prompt is **HOW** the loop survives a context reset). Propagation gate **CM-COVENANT-114-127-PROPAGATION** (literal 11.4.127 across the consumer fleet) + recommended gate **CM-HANDOFF-RESUMPTION-PROMPT-PRESENT** + paired §1.1 meta-test mutation (strip the literal → propagation gate **FAILS**; gate-code = separate work item).

Canonical authority: this Constitution.md §11.4.127 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--skip-handoff-prompt`, `--generic-prompt-OK`, `--no-resumption-sentence`, `--handoff-without-state` flag exists.

§11.4.128 — Always-on device-recording mandate (User mandate, 2026-06-06)

Forensic anchor — direct user mandate (2026-06-06): we MUST ALWAYS live-record all available data from all devices we use for testing or that we know manual testing is being performed on, done EXTRA carefully so the recording never harms the device, its performance, or causes side effects; recorded raw data MUST NOT be processed without need (token-conscious) and is ALWAYS excluded from version control + the code-intelligence index; only curated analysis/evidence is committed, and only during release preparation.

The mandate (ALL must hold). For EVERY test/debug device the project uses — and every device on which the operator is known to be performing manual testing — across EVERY reachable transport (USB or wireless ADB / SSH / serial / network introspection API), the project MUST ALWAYS live-record all available data sources it can later analyse: all activities, all logs (system log, kernel log, dropbox, crash/ANR traces), performance metrics (CPU / memory / I/O / thermal / load), every sink-side / downstream introspection report per §11.4.13, and any other live-changeable parameter. (1) **Extra-careful, side-effect-free.** Recording MUST be designed so it NEVER harms the device, degrades its performance, or causes a side effect — non-invasive read-only probes only, bounded sampling intervals, bounded write-volume, an observer-effect budget (no probe that mutates the system-under-test). A recorder that perturbs the device-under-test is a §11.4.128 violation, NOT evidence. (2) **Background + parallel + subagent-driven.** Recording MUST run in the background in parallel with the main work stream (per §11.4.103 continuous-parallel-stream), dispatched subagent-driven per §11.4.70 — never blocking the main stream, never the only thing the conductor does. (3) **Token-conscious — record-now, analyse-later.** Raw recorded data MUST NOT be processed / analysed without need; the default is capture-and-store. The ONLY standing exception is during release-tag preparation (the §11.4.40 full-suite retest / §11.4.42 release-gate) OR when the operator explicitly asks — at which point the relevant slice is analysed and the curated findings become captured-evidence. Analysing raw recordings continuously, eagerly, for no tracked reason burns tokens and is a §11.4.128 violation. (4) **Raw is git-ignored AND code-intelligence-excluded; only curated evidence is committed.** The raw recording corpus MUST be excluded from version control (with a §11.4.77 regeneration-mechanism declaration — the mechanism is “re-record by re-running the recorder against the device”, with the honest caveat that a past device-state is not byte-reproducible) AND excluded from the code-intelligence index per §11.4.78/§11.4.79 (a raw log/recording corpus is not source; indexing it pollutes the graph). Only the CURATED analysis / extracted evidence artefacts are committed — and only during release preparation per §11.4.83 (docs/qa/<run-id>/) — never the raw corpus. (5) **Deterministic archive layout.** The raw corpus MUST be laid out by date → combined-source-state-hash → device → recording-sequence so any artefact is traceable to the exact code state and device it came from: <recording-root>/YYYY-MM-DD/<combined main+submodules state hash>/<DEVICE>_<SERIAL>/recording_NNN/<files> (monotonic NNN per device per state-hash). (6) **Anti-bluff.** A recorder claimed running but producing no growing corpus is a §11.4 bluff; a curated finding cited as evidence MUST trace to a real raw-corpus path under the deterministic layout; the recorder’s own health (alive + corpus growing) is itself a captured-evidence assertion per §11.4.5/§11.4.69.

Classification: universal (§11.4.17) — a platform-neutral test-evidence-capture discipline reusable by ANY project with reachable test/debug devices; the consuming project supplies its concrete device serials, transport probes, recorder

helper, recording-root, and performance-metric set per §11.4.35. Composes with §11.4.2 (recorded-evidence — §11.4.128 guarantees the recording exists to BE the evidence) / §11.4.5 (captured-evidence quality) / §11.4.13 (sink-side introspection IS one recorded source) / §11.4.69 (sink-side positive-evidence taxonomy) / §11.4.40 + §11.4.42 (release-prep is the analyse-the-corpus trigger) / §11.4.70 (subagent-driven recorder dispatch) / §11.4.77 (raw corpus git-ignored WITH a regeneration-mechanism declaration) / §11.4.78 + §11.4.79 (raw corpus excluded from the code-intelligence index) / §11.4.83 (curated evidence lands under docs/qa/<run-id>/) / §11.4.103 (background parallel stream) / §11.4.119 (single-resource-owner — a read-only recorder never contends with an exclusive driver). Propagation gate CM-COVENANT-114-128-PROPAGATION (literal 11.4.128 across the consumer fleet) + recommended gate CM-DEVICE-RECORDING-ALWAYS-ON (recorder present + raw-corpus git-ignored + code-intelligence-excluded + deterministic layout) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: this Constitution.md §11.4.128 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no --skip-recording, --record-without-layout, --commit-raw-corpus, --index-raw-corpus, --analyse-corpus-always, --invasive-probe-OK flag exists.

§11.4.129 — Huge-blocker release protocol (User mandate, 2026-06-06)

Forensic anchor — direct user mandate (2026-06-06): when a huge blocker is discovered during release validation we MUST stop all testing, return to fixing ALL discovered issues, process all recorded data from the last session, land rock-solid fixes, author NEW validation+verification tests of ALL supported test types, rebuild, reflash, and RESTART the full validation+verification of every fix/change from the last release tag to now — on both devices in parallel, recorded, with real physical captured proofs and no bluff.

When a HUGE BLOCKER (a release-blocking-severity defect — a core user-facing capability broken, a regression that invalidates the in-flight validation cycle, or a defect whose blast radius reaches the batch's other fixes) is discovered during release validation, the agent MUST execute the following protocol in order, with NO partial / spot-check shortcut: (1) **STOP all testing** immediately on every device (the §11.4.4 test-interrupt-on-discovery STOP, escalated to release granularity — a validation cycle that continues past a huge blocker produces “all green” summaries on a known-broken build, the exact §11.4 PASS-bluff). (2) **Return to fixing ALL discovered issues** — not only the huge blocker but every defect surfaced this cycle; investigate each to root cause per §11.4.102 systematic-debugging, isolate regressions against the last known-good tag per §11.4.114. (3) **Process all recorded data from the last session** — analyse the §11.4.128 raw-corpus slice for the cycle (this is the §11.4.128(3) release-prep analyse-trigger) so no defect the recordings already captured is missed. (4) **Land rock-solid fixes** — per §11.4.123 (rock-solid proof or deep research), §11.4.43/§11.4.115 (RED-on-broken-artifact then GREEN), §11.4.9 batch-source-fixes-before-rebuild. (5) **Author NEW validation+verification tests of ALL supported test types** — per §11.4.27 (unit / integration / e2e / full-automation / security / stress / chaos / performance / UI / UX / Challenges / HelixQA) + §11.4.85 stress+chaos, each anti-bluff with captured

evidence + paired §1.1 mutation. (6) **Rebuild + reflash** — full rebuild (NOT module-only) for retest-baseline integrity, deploy to a CLEAN target per §11.4.108 (no stale-overlay shadow). (7) **RESTART the full validation+verification from the last release tag to now** — re-validate EVERY fix/change landed since the last tag (per §11.4.40 full-suite-retest-before-tag — the huge-blocker discovery invalidated the in-flight cycle, so the cycle restarts from the beginning, NOT resumes), running on both/all owned devices IN PARALLEL per §11.4.103/ §11.4.119 (single-resource-owner partitioning), every device's run RECORDED per §11.4.128, producing REAL physical captured proofs per §11.4.5/§11.4.69/ §11.4.107 with no bluff of any kind. This anchor BINDS the existing release-cycle anchors for the huge-blocker case — it adds the STOP→fix-all→process-recordings→new-tests-all-types→rebuild→reflash→full-restart sequencing and the “restart, never resume” rule, citing §11.4.4 + §11.4.40 + §11.4.108 + §11.4.114/ §11.4.115 + §11.4.123 + §11.4.128 rather than duplicating their text.

Classification: universal (§11.4.17) — a platform-neutral release-discipline reusable by ANY project that runs a release-validation cycle; the consuming project supplies its concrete device set, build/flash mechanism, and last-known-good tag per §11.4.35. Composes with §11.4.4 (test-interrupt-on-discovery — §11.4.129 is its release-granularity escalation) / §11.4.40 (full-suite retest before tag — §11.4.129 mandates the RESTART, not resume) / §11.4.42 (iteration discipline) / §11.4.9 (batch-source-fixes-before-rebuild) / §11.4.27 (all-test-types coverage) / §11.4.85 (stress+chaos) / §11.4.102 (systematic-debugging) / §11.4.108 (clean-target rebuild+reflash+runtime-signature) / §11.4.114 + §11.4.115 (regression isolation + RED-on-broken-artifact) / §11.4.123 (rock-solid proof) / §11.4.128 (process the recorded corpus) / §11.4.103 + §11.4.119 (parallel recorded multi-device). Propagation gate CM-COVENANT-114-129-PROPAGATION (literal 11.4.129 across the consumer fleet) + recommended gate CM-HUGE-BLOCKER-FULL-RESTART (a release-blocker discovery during validation triggers a from-the-last-tag full restart, not a resume) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: this Constitution.md §11.4.129 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no --resume-after-blocker, --spot-validate-after-fix, --skip-recording-analysis, --skip-new-tests, --module-only-after-blocker, --single-device-restart flag exists.

§11.4.130 — Post-remediation validate-the-fix-FIRST-after-redeploy (User mandate, 2026-06-06)

Forensic anchor — direct user mandate (2026-06-06): when a blocker discovered during release validation is fixed and a new artifact is produced + redeployed, we MUST first re-test the SPECIFIC last-failing features and validate+verify the just-incorporated fixes BEFORE proceeding to the broader / full validation — a first fix attempt may not work, may be incomplete, or may regress again, so confirming the targeted fix on the new artifact first prevents wasting a full multi-hour cycle on a still-broken fix.

When a blocker / critical failure is discovered during release validation and then FIXED, and a new artifact is produced (rebuild / new flashing image / redeploy) and the target is reflashed / redistributed / updated, the agent MUST execute the

following ordering with NO shortcut: (1) **Re-test the SPECIFIC last-failing features FIRST** — run the targeted guard tests for exactly the defects this fix addressed, BEFORE any broader / full-suite validation. (2) **Validate + verify the just-incorporated fixes with real captured evidence** — the RED test that reproduced the defect on the pre-fix artifact (per §11.4.115) MUST now flip GREEN at RED_MODE=0 on the new artifact, AND the fix’s declared runtime signature (per §11.4.108) MUST verify on the CLEAN target the redeploy produced. Metadata-only / config-only / absence-of-error / grep-without-runtime PASS are forbidden (§11.4 / §11.4.1) — the proof is captured per §11.4.5 / §11.4.69 / §11.4.107 / §11.4.123. (3) **Only after the targeted fix is CONFIRMED working on the new artifact** does the agent proceed to the §11.4.40 full retest of every fix/ change from the last release tag to now.

Rationale. A first fix attempt may not work, may be incomplete, or may regress again under the new artifact. Confirming the targeted fix FIRST catches a fix-did-not-take case immediately — instead of discovering it hours later at the END of a full multi-hour validation cycle (and then having to restart the whole cycle per §11.4.129). The targeted re-test is cheap; the full cycle is expensive; ordering cheap-confirmation-first is the §11.4.82 iteration-speedup discipline applied to the post-blocker reflash. This is the §11.4.46 recent-work-validation gate specialised for the post-blocker-reflash case: the §11.4.129 huge-blocker protocol’s step (7) full-restart MUST be preceded by this targeted-fix-first confirmation phase.

Honest boundary (§11.4.6): “the fix probably took” is not “the fix took” — the targeted RED→GREEN flip + runtime-signature verification on the new artifact is the captured proof; proceeding to the full cycle without it is a §11.4.130 violation. A targeted re-test that still FAILs on the new artifact is a finding — return to the §11.4.114/§11.4.115 isolate→RED→fix loop and produce another artifact; do NOT proceed to the full cycle on a still-broken fix.

Classification: universal (§11.4.17) — a platform-neutral release-discipline reusable by ANY project that fixes a blocker, redeploys, and re-validates; the consuming project supplies its concrete targeted-guard-test set, build/flash mechanism, and clean-deployment / runtime-signature observables per §11.4.35. Composes with §11.4.4 (test-interrupt-on-discovery — the blocker that triggered the fix) / §11.4.40 (full-suite retest before tag — §11.4.130 is the targeted-first phase that GATES the full retest) / §11.4.46 (validate-recent-work-before-post-flash-tests — §11.4.130 is its post-blocker-reflash specialisation) / §11.4.108 (runtime-signature-as-definition-of-done on the clean target the redeploy produced) / §11.4.114 (regression isolation — a still-failing targeted re-test re-enters the isolation loop) / §11.4.115 (RED-baseline + polarity-switch — the RED test flips GREEN on the new artifact) / §11.4.123 (rock-solid proof — captured evidence, never “probably took”) / §11.4.129 (huge-blocker protocol — §11.4.130 is its post-reflash targeted-confirmation phase that precedes the step-7 full-restart) / §11.4.82 (iteration-speedup — cheap-confirmation-first). Propagation gate CM-COVENANT-114-130-PROPAGATION (literal 11.4.130 across the consumer fleet) + recommended gate CM-FIX-FIRST-AFTER-REDEPLOY (after a post-blocker redeploy, the targeted last-failing-feature guard tests run + pass on the new artifact before the full retest starts) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: this Constitution.md §11.4.130 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--skip-targeted-retest`, `--full-cycle-first`, `--assume-fix-took`, `--validate-fix-at-end`, `--skip-red-green-flip-on-new-artifact` flag exists.

§11.4.131 — Standing session-resumption file mandate (User mandate, 2026-06-07)

Forensic anchor — verbatim user mandate (2026-06-07): “Make this markdown a standard file which will be written EVERY TIME when we need fresh session out of the box! It MUST BE always up to date and in sync so whenever new session is created all we have to do is just point to it!”

Every project MUST maintain a SINGLE canonical, always-current **session-resumption file** at a fixed, project-declared standard path (the project declares the exact path once per §11.4.35 and never moves it without a §11.4.66 operator decision). This file is the OUT-OF-THE-BOX entry point for any fresh session: creating a new session requires ONLY pointing the new agent at this one file. §11.4.131 promotes §11.4.127 (which mandated PREPARING a resumption prompt on demand) into a STANDING, version-controlled ARTIFACT — the resumption prompt is no longer an ephemeral in-message reply but a persistent file that is ALWAYS present and ALWAYS in sync.

1. **Existence + fixed path.** The file MUST exist at the declared canonical path at all times; its location is encoded as a literal path string in the project-layer instantiation (§11.4.35) and never silently moved.
2. **Always written + always synced.** The file MUST be (re)written whenever a fresh session is needed OR whenever the live project state materially changes (new HEAD, build/artifact id, phase transition, device/target state, in-flight job, blocking decision) — the §12.10 CONTINUATION.md trigger set applied to the resumption file. A stale resumption file (describing a prior HEAD / phase / artifact) is a §11.4.131 violation of the same severity class as a §12.10 stale-CONTINUATION violation: it silently breaks the zero-loss-resumption guarantee.
3. **Content (composes §11.4.127).** The file MUST carry both the §11.4.127 SHORT (one-paste first-sentence) AND FULL (paste-ready block) variants; point to the live handoff docs read FIRST (`.remember/remember.md` if present + `docs/CONTINUATION.md` per §12.10) + `git fetch`; embed the exact live-state anchors (HEAD commit, build/artifact ids + checksums, device/target serials, in-flight PIDs + log paths, captured-evidence paths); state the current PHASE + immediate NEXT action + terminal goal; and restate the binding constraints (anti-bluff §11.4, no-force-push §11.4.113, exact version/naming, hardware/target gotchas). It MUST be MOMENT-VALID, never a generic template (§11.4.6).
4. **Export + freshness.** The file is in §11.4.65 scope → synchronized `.html` / `.pdf` siblings refreshed on every update; a §11.4.44 revision header tracks its monotonic revision + last-modified timestamp.
5. **Out-of-the-box resumption.** A fresh session, given ONLY this file’s path, MUST be able to fully resume with zero additional context. The file IS the session bootstrap.

Composes §12.10 (CONTINUATION.md live-state — §11.4.131 is the dedicated paste-ready-prompt sibling of it) / §11.4.127 (resumption-prompt discipline promoted to a standing file) / §11.4.65 (universal export) / §11.4.44 (revision header) / §11.4.6 (moment-valid, never generic) / §11.4.66 (path change is an operator decision) / §11.4.126 (default-loop continuity across sessions). Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-131-PROPAGATION (literal 11.4.131 across the consumer fleet) + recommended gate CM-SESSION-RESUMPTION-FILE-PRESENT (the declared file exists + carries SHORT+FULL + revision header + is fresh vs HEAD) + paired §1.1 meta-test mutation (strip the file / make it stale → gate FAILs; gate-code = separate work item).

Canonical authority: this Constitution.md §11.4.131 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no --skip-resumption-file, --ephemeral-prompt-only, --stale-resumption-OK, --generic-template-OK flag exists.

§11.4.132 — Risk-ordered validation priority mandate (User mandate, 2026-06-07)

Forensic anchor — verbatim user mandate (2026-06-07): “We MUST ALWAYS first test and validate features, functionalities and fixes/changes that have been worked most recently, the ones which were most problematic, which have the most chance to crash or break again, the ones which have been re-opened the most times! Then, after we validate and verify all this with real (physical) proofs and hard evidence, with no false results and bluffs of any kind, we continue with all other existing tests in the test suites! This IS MANDATORY.”

Tests, validations, and verifications MUST run in **RISK-DESCENDING order** — the highest-risk set FIRST, and only AFTER that set is fully GREEN with real (physical) captured evidence does the remainder of the suite run. The risk ranking is computed from a closed set of factors, highest-risk first:

1. **Most-recently-worked** — features / functionalities / fixes / changes touched most recently (the freshest code is the least-exercised in the wild and the most likely to carry an undiscovered defect).
2. **Historically most-problematic** — the items with the longest defect history, the most prior fixes, the most prior failures.
3. **Highest crash/break/regress likelihood** — items whose blast radius, complexity, or dependency surface gives them the greatest chance to crash or break again.
4. **Most-reopened** — items with the highest reopen count per §11.4.55 (a high reopens-count is the strongest empirical signal of fragility — a defect that keeps coming back).

The highest-risk set, ordered by (a)–(d), is validated FIRST. Each item in that set MUST pass with real (physical) captured evidence per §11.4.5 / §11.4.69 / §11.4.107 — no metadata-only / config-only / absence-of-error / grep-without-runtime PASS (§11.4 / §11.4.1), no false results, no bluff of any kind (§11.4.6). ONLY AFTER the

entire highest-risk set is GREEN with captured proof does the agent proceed to all other existing tests in the test suites. Running the full suite in arbitrary order, or running lower-risk tests before the highest-risk set is GREEN, is a §11.4.132 violation: it spends the operator's scarce validation window on the items least likely to fail while deferring the items most likely to fail.

§11.4.132 REFINES / STRENGTHENS §11.4.130 (post-remediation validate-the-fix-FIRST-after-redeploy — §11.4.130 says re-test the just-fixed items first; §11.4.132 generalises “first” to the full risk-ordered set) + §11.4.46 (validate-recent-work-before-post-flash-tests — §11.4.46 gates the full suite on recent-work GREEN; §11.4.132 adds the explicit risk-ordering WITHIN the recent / high-risk set) + §11.4.42 (iteration-discipline priority order — §11.4.132 applies the same priority discipline to VALIDATION ordering, not just implementation ordering). The risk ranking is the validation-layer analogue of the §11.4.42 / §11.4.72 implementation-layer priority queue.

Classification: universal (§11.4.17) — a platform-neutral validation-ordering discipline reusable by ANY project; the consuming project supplies its concrete recency / problematic-history / reopen-count sources (e.g. the §11.4.93 workable-items DB reopens_count + last_modified columns) per §11.4.35. Composes §11.4.4 / §11.4.5 / §11.4.6 / §11.4.7 / §11.4.40 / §11.4.42 / §11.4.46 / §11.4.50 / §11.4.55 / §11.4.69 / §11.4.107 / §11.4.130. Propagation gate CM-COVENANT-114-132-PROPAGATION (literal 11.4.132) + recommended gate CM-RISK-ORDERED-VALIDATION-PRIORITY + paired §1.1 meta-test mutation (gate-code = separate work item).

Canonical authority: this Constitution.md §11.4.132 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no --skip-risk-ordering, --any-order-OK, --suite-order-fixed flag exists.

§11.4.133 — Target-System + hardware safety mandate (User mandate, 2026-06-08)

Forensic anchor — verbatim user mandate (2026-06-08): “Make sure that all changes we do to the System are ALWAYS safe for the System itself and for the hardware the system runs on! This is MANDATORY.”

Every change to the TARGET system — firmware, kernel, init / boot scripts, drivers, sysfs / devfreq / voltage / clock / thermal / regulator register writes, partition layout / bootloader / U-Boot, HAL, framework, prebuilts, device configuration, or any equivalent on a non-Android target — MUST ALWAYS be safe for BOTH: (a) **the target System itself** — the change MUST NOT brick the device, induce a boot-loop, corrupt user / system data, or render the device unrecoverable; AND (b) **the hardware the System runs on** — the change MUST NOT drive any hardware-control surface beyond its safe electrical / thermal / voltage / clock envelope, and MUST NOT damage panels, storage, radios, regulators, or any other physical component.

Concrete obligations (ALL hold):

1. **Reversible-first.** Prefer reversible changes; when a change is irreversible or high-blast-radius (bootloader / U-Boot / partition layout / one-way fuse / OTP), verify it against a known-good value (MD5 / checksum / golden manifest) BEFORE applying, and capture a pre-op backup of any state the change overwrites (per §9.2). A wrong irreversible write that bricks the device is the worst-case the mandate exists to prevent.
2. **No unverified hardware-control writes.** NEVER write an unverified value to a hardware-control sysfs node / register / firmware field that governs voltage, clock frequency, regulator output, thermal-throttle threshold, current limit, or any other physical-safety parameter. Every such value MUST be within the component's datasheet / vendor-documented safe range; the safe range MUST be established as FACT (cited datasheet / vendor reference per §11.4.8 + §11.4.99), never guessed (§11.4.6). Pinning a max OPP / disabling a throttle / raising a regulator output without proving the cooling design and electrical envelope tolerate it is forbidden.
3. **Thermal / performance changes respect the cooling design.** Any change that raises sustained power, frequency, or voltage (e.g. forcing a performance governor, pinning the top OPP, disabling idle / thermal management) MUST be validated against the device's actual cooling capacity — captured thermal evidence (sustained-load temperature $\leq$ the documented limit) is the proof; "it boots" is not proof it is thermally safe.
4. **Flashing uses the sanctioned tool + verified image.** Deploying firmware MUST use the project's sanctioned flashing tool against a freshly-built, integrity-verified image — never an ad-hoc partition write, never a stale / unverified artifact, never a tool / image whose provenance is unestablished.
5. **Unprovable-safety $\Rightarrow$ blocked.** A change whose safety for the target System and its hardware cannot be established from captured evidence MUST be treated as UNSAFE and blocked — proceed only after the safe outcome is determined as FACT (§11.4.6), the change is reversible OR backed up (§11.4.101 reversible-first + §9.2), and (when the safe choice still cannot be determined) the operator is asked via §11.4.66 / §11.4.101's block-only-when rule. "Probably safe" is the exact guess this mandate forbids.

Distinct from §12 host-session safety. §12 (and §12.6 / §12.7 / §12.8 / §12.9 / §12.10) protects the DEVELOPER's HOST machine and session (no host suspend / logout / OOM cascade / unbounded-memory build). §11.4.133 protects the TARGET device the System is built FOR and runs ON. Both apply; neither weakens the other. A build step can be host-safe (bounded memory, no host suspend) yet still ship a target-unsafe change (an over-voltage regulator write) — §11.4.133 is the gate for the latter.

Classification: universal (§11.4.17) — a platform-neutral target-safety discipline reusable by ANY project that produces firmware / kernel / driver / device-config changes deployed onto target hardware; the consuming project supplies its concrete hardware-control surfaces, datasheet-safe ranges, known-good bootloader / image hashes, and sanctioned flashing tool per §11.4.35. Composes §12 (host-session safety — sibling-not-overlapping protection scope) / §11.4.6 (no-guessing — the safe range / known-good value is FACT, never assumed) / §11.4.101 (reversible-first + autonomous-decision-over-blocking — irreversible high-blast-radius hardware

writes are the class that escalates to an operator question) / §11.4.108 (the change's runtime signature on a clean target proves it landed AND is safe) / §11.4.123 (rock-solid proof — safety claims need captured evidence, never a metadata-only assertion). Propagation gate CM-COVENANT-114-133-PROPAGATION (literal 11.4.133 across the consumer fleet) + recommended gate CM-TARGET-HARDWARE-SAFETY + paired §1.1 meta-test mutation (gate-code = separate work item).

Canonical authority: this Constitution.md §11.4.133 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no --unsafe-hardware-write, --skip-system-safety, --brick-risk-accepted flag exists.

§11.4.134 — Code-review iterate-until-GO + rock-solid-evidence mandate (User mandate, 2026-06-08)

Forensic anchor — verbatim user mandate (2026-06-08): “For any fixes/changes given back to us for re-work by the code-review process, once we fix/improve everything per the code-review's requests, we MUST RE-RUN code-review AGAIN until we get a GO from it with NO new issues reported or warnings of any kind! All results produced by this whole process MUST ALWAYS give us rock-solid PHYSICAL evidence that the fixed/improved codebase really works now as expected, with no false results and no bluff(s) of any kind.”

When the §11.4.125 code-review returns ANY finding — BLOCKING, nit, or warning — and the author fixes/improves the batch per that review, the code review MUST BE RE-RUN, and MUST KEEP being re-run after each remediation round, until it returns a clean GO with ZERO new issues AND ZERO warnings of any kind. A single pass that “addressed the findings” is NOT sufficient: the corrected batch MUST pass a FRESH adversarial review (a re-review can surface NEW findings introduced by the very fixes that closed the prior ones — the §11.4.1 fix-A-creates-B failure mode). The loop terminates ONLY on a clean GO (no new findings, no warnings); a residual warning is itself a finding that re-arms the loop.

Every round's verdict AND every fix's validation MUST carry rock-solid PHYSICAL captured evidence per §11.4.5 / §11.4.69 / §11.4.107 (captured audio / video / sysfs / dumphsys / sink-side / runtime-signature) proving the fixed/improved codebase REALLY works as expected on the target — never metadata-only / configuration-only / absence-of-error / grep-without-runtime. No false results, no bluff of any kind, at any round. A reported GO unbacked by captured physical evidence is itself a §11.4 PASS-bluff at the review-loop layer.

§11.4.134 REFINES / STRENGTHENS §11.4.125 (which requires the review to iterate “until no blocking findings remain”): it makes the loop EXPLICIT (re-run after every remediation round, not once), raises the termination bar to ZERO findings AND ZERO warnings (not merely zero-blocking), and BINDS rock-solid physical evidence to every round (not just the final one). Classification: universal (§11.4.17). Composes §11.4.125 / §11.4.1 / §11.4.4 / §11.4.5 / §11.4.6 / §11.4.69 / §11.4.107 / §11.4.50 / §11.4.108 / §11.4.123. Propagation gate CM-COVENANT-114-134-PROPAGATION (literal 11.4.134) + recommended gate CM-CODE-REVIEW-ITERATE-UNTIL-GO + paired §1.1 meta-test mutation (gate-code = separate work item).

Canonical authority: this Constitution.md §11.4.134 in the HelixConstitution submodule. All consuming projects restate + cite via §11.4.35 inheritance.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--skip-rereview`, `--single-review-pass`, `--warnings-ok`, `--evidence-optional` flag exists.

§11.4.135 — Standing regression-guard suite + every-fixed-defect-gets-a-permanent-regression-test (User mandate, 2026-06-08)

Forensic anchor (the canonical case this anchor exists to prevent, FACT): the wrong-subtitle-on-2nd-display defect was “fixed” in a prior release via a source-side `CONTROL_MENU_LABEL_DENYLIST` that NO test mirrored or re-ran, so the NEXT chrome class (a settings-menu label) recurred silently while the GREEN suite passed.

Every project MUST maintain a STANDING regression-guard suite that runs on EVERY build+deploy and BLOCKS the release tag on any failure. Every closed defect (the project’s stable ticket id, e.g. ATM-NNN) MUST, in the SAME commit as its fix (extending the §11.4.43 DOCUMENT step), register a permanent §11.4.115 RED-on-broken-artifact regression test into the suite — `RED_MODE=1` capturing the historical defect on a pre-fix artifact (the proof the guard is real), `RED_MODE=0` the standing GREEN guard asserting the defect is ABSENT. A closure without a registered guard is a §11.4.123 violation (closure without permanent regression proof). The suite runs FIRST in the post-deploy cycle (highest-risk set per §11.4.132 risk-ordering) and is a §11.4.40 release-gate blocker. This is industry-standard bug-driven testing (cite: Google for Developers content-driven testing <https://developers.google.com/solutions/content-driven/backend/testing>; AOSP CTS/Tradefed continuous-regression model <https://source.android.com/docs/compatibility/cts>) made mechanical + enforced.

Classification: universal (§11.4.17). Composes §11.4.4 / §11.4.40 / §11.4.43 / §11.4.46 / §11.4.50 / §11.4.107 / §11.4.108 / §11.4.115 / §11.4.118 / §11.4.123 / §11.4.124 / §11.4.130 / §11.4.132. Recommended gates `CM-REGRESSION-GUARD-REGISTERED` + `CM-REGRESSION-GUARD-SUITE-WIRED` + `CM-COVENANT-114-135-PROPAGATION` (literal 11.4.135) + paired §1.1 meta-test mutations (gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md §11.4.135. Non-compliance is a release blocker. No escape hatch — no `--skip-regression-guard`, `--no-guard-on-close`, `--guard-optional` flag.

§11.4.136 — Real-content end-to-end playback-test mandate (User mandate, 2026-06-08)

Refines/strengthens §11.4.107. Any test asserting media playback works MUST drive REAL content (catalog stream or offline reference clip) through the user’s path (§11.4.48 UI-driven → §11.4.117 CV/OCR fallback) and assert it genuinely PLAYS via the §11.4.107 liveness battery PLUS a decoder-health census — a numeric drop-buffer budget, no buffer-timestamp re-order/discard, no codec-reject (cite: Android/Media3 ExoPlayer OEM pre-OTA playback-test mandate <https://developer.android.com/media/media3/exoplayer/oems> — “too many dropped buffers” >25, “unexpected presentation timestamp”, “test timed out” failure modes).

Metadata-only / launch-only / registration-only / single-frame / config-only PASS is forbidden (§11.4 / §11.4.1). A golden/reference clip corpus (cite: BBC ExoPlayer testing samples <https://github.com/bbc/exoplayer-testing-samples>; USPTO 8,387,015 empirical media-player testing) is the offline ground-truth.

Classification: universal (§11.4.17). Composes §11.4.5 / §11.4.48 / §11.4.50 / §11.4.107 / §11.4.117 / §11.4.123 / §11.4.13 / §11.4.69. Recommended gate CM-REAL-CONTENT-PLAYBACK-TEST + CM-COVENANT-114-136-PROPAGATION (literal 11.4.136) + paired §1.1 meta-test mutation (gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.136. Non-compliance is a release blocker. No escape hatch — no `--launch-proves-playback`, `--skip-decoder-health`, `--metadata-playback-pass-suffices` flag.

§11.4.137 — Subtitle/caption content-correctness oracle + secure-display-proxy-honesty mandate (User mandate, 2026-06-08)

Refines §11.4.117 (pixel-oracle) + §11.4.107 (liveness) + §11.4.112 (structural-impossibility).

Forensic anchor (FACT): tests tasked to “physically verify the 2nd-display (or 1st where no TV) subtitle” PASSEd GREEN while subtitles did NOT show / showed WRONG, because (a) the streaming app’s player surface is FLAG_SECURE so `screencap -d <secondary>` returns BLACK (cite: Android Developers <https://developer.android.com/security/fraud-prevention/activities>; Nightwatch <https://www.nightwatchcybersecurity.com/2016/04/13/research-securing-android-applications-from-screen-capture/>; WithSecure <https://labs.withsecure.com/advisories/screencapture-via-ui-overlays-in-mediaprojection>) → autonomous PIXEL verification is structurally impossible (§11.4.112), so the test fell back to a non-pixel proxy (the accessibility-scraped forwarded cue / `persist.atmosphere.subdebug logcat`), and (b) the proxy’s validation accepted a chrome/menu LABEL (Аудио и субтитры) as a valid subtitle because the prose floor accepted any multibyte prose and NO menu-label denylist + NO position/cadence check existed.

The mandate: a test asserting a subtitle/caption is correct MUST classify the *content class* of the rendered/forwarded cue — a present cue is NOT a correct cue. A cue is CHROME (FAIL) if it is a known control/menu label (a closed multilingual deny-list MIRRORED from the source-side denylist, case-folded both sides incl. non-ASCII scripts), is time-counter/numeric chrome, is NOT prose, is OUTSIDE the lower safe-title position band where caption position is available (cite CEA-708 9-anchor caption-window grid: captions live in the lower-center safe band, chrome elsewhere <https://www.gte-media.com/cea-608-and-cea-708-captions/>), OR is STATIC across the window (a real subtitle changes → dialogue cadence; a menu label does not → ≥2 distinct prose cues required, a metamorphic relation per https://en.wikipedia.org/wiki/Metamorphic_testing). A cue is DIALOGUE (PASS) only when prose, not-denied, not-chrome, position-ok, AND cadence ≥2 OR it fuzzy-matches the SOURCE-extracted expected cue via normalized edit distance (§11.4.123 host-side ground truth; cite SubER <https://github.com/apptek/SubER>). The content oracle MUST be self-validated golden-good/golden-bad (§11.4.107(10)) and the test-side deny-list MUST be verified present in the SHIPPED artifact (§11.4.108) — a source-green denylist with no test mirror + no artifact check is the exact recurrence pattern this anchor forbids.

Secure-display-proxy honesty (§11.4.112): where FLAG_SECURE makes pixel verification structurally impossible, the rock-solid autonomous proof is the player's own caption telemetry (the subdebug/Cue ground-truth channel) + source-track presence + content-class oracle — NEVER a faked pixel “physical” pass; the human-eye pixel confirmation is operator_attended (§11.4.52) with a tracked migration item. App-agnostic (keys off content class, not per-app rule) so it is the standing guard for the whole roster.

Classification: universal (§11.4.17). Composes §11.4.3 / §11.4.5 / §11.4.6 / §11.4.107 / §11.4.108 / §11.4.112 / §11.4.115 / §11.4.117 / §11.4.123 / §11.4.13 / §11.4.69. Recommended gate CM-SUBTITLE-CONTENT-CORRECTNESS-ORACLE + CM-COVENANT-114-137-PROPAGATION (literal 11.4.137) + paired §1.1 meta-test mutation (strip the denylist/position/cadence check → golden-bad Аудио и субтитры PASSES → gate FAILs).

Canonical authority: constitution submodule Constitution.md §11.4.137. Non-compliance is a release blocker. No escape hatch — no --present-cue-is-correct, --skip-chrome-oracle, --length-heuristic-suffices, --pixel-pass-on-secure-display, --skip-position-check, --skip-cadence-check flag.

§11.4.138 — Operator-escape => mandatory bluff-audit + permanent guard (User mandate, 2026-06-08)

When the operator (or any out-of-band channel) finds a defect that the GREEN test suite passed, this is by definition a §11.4 PASS-bluff — it MUST trigger, before the fix is closed: (1) a §11.4.102 systematic-debugging pass to FACT-root-cause; (2) a bluff-audit identifying the EXACT assertion that should have caught it but didn't, cited to file:line (canonical example: lib/subtitle_content_validation.sh:sub_is_prose() returning TRUE for Аудио и субтитры); (3) a permanent §11.4.135 regression guard registered in the SAME commit as the fix, with its §11.4.115 RED capturing the operator-found defect; (4) the bluff-audit committed under docs/research/<scope>/<defect>_bluff_audit/. Closing an operator-found defect WITHOUT the bluff-audit + permanent guard is itself a §11.4 violation (the bluff that let it through is still live and the defect will recur — the operator's “repeatedly reworking + reopening the same issues with no real progress” pain).

Classification: universal (§11.4.17). Composes §11.4 / §11.4.1 / §11.4.102 / §11.4.108 / §11.4.115 / §11.4.118 / §11.4.123 / §11.4.135. Recommended gate CM-OPERATOR-ESCAPE-BLUFF-AUDIT + CM-COVENANT-114-138-PROPAGATION (literal 11.4.138) + paired §1.1 meta-test mutation (gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md §11.4.138. Non-compliance is a release blocker. No escape hatch — no --close-without-bluff-audit, --operator-find-is-just-a-bug, --skip-permanent-guard flag.

§11.4.139 — Fresh-process clean-artifact runtime-signature mandate (User mandate, 2026-06-08)

Refines §11.4.108. Before any post-deploy validation — ESPECIALLY a non-pixel proxy verification (the subdebug/accessibility-cue channel used for FLAG_SECURE displays) — the harness MUST assert running-artifact == built-

artifact: the deploy yielded a CLEAN target (mutable-overlay/userdata wiped) OR a pre-validation check proves no stale overlay shadows the deployed code (e.g. every guarded package — incl. the Presenter that emits the subtitle cue — resolves to the system partition, no per-user override). A stale shadow of the cue-emitting component (e.g. a Presenter APK predating the denylist) makes the proxy report on code that was never deployed — any PASS is a §11.4 PASS-bluff. Each fix declares ONE machine-checkable runtime signature verified on the clean target (the §11.4.108 registry IS the definition of done); for the subtitle class the signature is “the shipped Presenter APK contains the denylist literal (case-insensitive) AND the subdebug channel emits candidate REJECTED reason=chrome-label for a menu label.”

Classification: universal (§11.4.17). Composes §11.4.46 / §11.4.108 / §11.4.130 / §11.4.135 / §11.4.137. Recommended gate CM-CLEAN-ARTIFACT-RUNTIME-SIGNATURE + CM-COVENANT-114-139-PROPAGATION (literal 11.4.139) + paired §1.1 meta-test mutation (gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md §11.4.139. Non-compliance is a release blocker. No escape hatch — no --validate-against-running-state, --skip-clean-precondition, --shadow-OK flag.

§11.4.140 — Universal action-prefix system (ACTION_NAME ::) (User mandate, 2026-06-09)

Forensic anchor — verbatim user mandate (2026-06-09):

“When a user prompt STARTS with ACTION_NAME :: (e.g. BACKGROUND :: IMPORTANT: do X...) the agent MUST replace the ACTION_NAME :: prefix with that action’s registered expansion text, then execute the rest. For BACKGROUND :: the expansion is: ‘The following prompt that we will provide MUST BE executed in background in parallel with all main work streams using the subagents-driven development approach! All work done MUST PRODUCE rock solid evidence covered with hard physical proof(s) that all done is working as expected and as specified without any false results and without any bluff!’. The mechanism MUST be UNIVERSAL/extensible — more actions will be added later, each with its own expansion + rules. It MUST work with EVERY CLI agent, be fully decoupled + reusable by every project that includes the constitution submodule, and load + execute out of the box.”

Every project under this Constitution MUST support a universal, extensible **action-prefix system**: when a user prompt’s FIRST non-blank line starts with an uppercase action token followed by " :: " (grammar $\wedge([A-Z][A-Z0-9_]*)::$ — anchored at line start, UPPERCASE-only token, exactly one space on each side of ::), the agent MUST (1) look the token up in the shared **action registry** (tracked data file constitution/actions/registry.yaml, or \$HELIX_ACTION_REGISTRY); (2) if the token is a registered action, REPLACE the ACTION_NAME :: prefix with that action’s registered expansion text and apply its rules; (3) execute the REMAINDER of the prompt under the expanded instruction. The system is the C-preprocessor expand-then-rescan model applied to prompts (a line-anchored single-token macro): detect → substitute the registered expansion → execute the residual.

Two-layer architecture (both mandatory; LAYER 1 is the universal floor).

(LAYER 1 — universal, always-on, out-of-the-box) the action-prefix recognition instruction MUST be mirrored into EVERY agent context carrier the constitution maintains (CLAUDE.md, AGENTS.md, QWEN.md, GEMINI.md per §11.4.35) so the agent itself recognises and applies the prefix on every CLI agent — Claude Code, Gemini CLI, Qwen Code, OpenAI Codex CLI, GitHub Copilot CLI, Cursor, Aider, Cline, Continue, Roo Code, and any future agent that reads one of those carriers. (LAYER 2 — mechanical, where the agent exposes a pre-submit seam) a prompt-preprocessing hook reading the SAME registry applies the expansion deterministically (Claude Code UserPromptSubmit / UserPromptExpansion hook injecting the expansion via additionalContext; per-agent slash-command equivalents generated from the registry for Gemini/Qwen/Codex). LAYER 2 is the §11.4.109 anti-forgetting upgrade applied to prompt prefixes — the expansion holds even if model recall lapses; LAYER 1 guarantees the system works everywhere with zero extra setup. Honest §11.4.3 boundary (§11.4.6): transparent mechanical free-form ^PREFIX :: interception is genuinely available ONLY on Claude Code today; on every other agent the free-form form is honoured by LAYER 1 (the agent self-applies) and the slash-command equivalent is the mechanical convenience — this split is documented, never papered over.

Grammar (mandatory, all hold). (1) Anchored at the start of the first non-blank line ONLY; mid-prose tokens never match. (2) UPPERCASE-only token [A-Z] [A-Z0-9_]*; lowercase never matches. (3) Separator is exactly " :: " (space-colon-colon-space) — avoids C++ Foo::Bar, YAML key: value, URLs. (4) Stacked prefixes (A :: B :: rest) apply outer-to-inner, left-to-right (expand A, rescan, expand B, then the residual is the task). (5) A leading \ escapes the prefix (literal, no expansion) so action names can be discussed. (6) An unknown token that matches the grammar shape but is NOT registered is NEVER silently expanded or silently dropped — the agent asks which registered action was meant (§11.4.66 + §11.4.105) or treats it literally; it NEVER invents an expansion (§11.4.6). (7) Any prompt not satisfying the grammar is an ordinary prompt; the system is a no-op. The carrier-block GRAMMAR_ADDENDUM (2026-06-09) additionally recognises three EQUIVALENT forms of every action — PREFIX::ACTION_NAME :: <rest> (namespaced ::), /ACTION_NAME <rest> (bare slash), /PREFIX::ACTION_NAME <rest> (namespaced slash) — where PREFIX is an action NAMESPACE (reserved default DEFAULT) and the namespace separator :: carries NO surrounding spaces (distinct from the action-body " :: "); forms 3/4 are honoured with the same expansion/execution, the bare slash form yielding to a colliding built-in/host slash command while the namespaced slash form is always unambiguous.

Registry is the single source of truth + the extension contract. Adding a new action = adding ONE actions[] row (name, version, summary ≥6 words per §11.4.91, expansion, optional rules + composes_with) — no code change in either layer (the registry is data, not code). Both layers read the same file (§11.4.93-style single-source-of-truth). The first registry entry is BACKGROUND with the operator's verbatim expansion above; it composes §11.4.20/§11.4.70 (subagent-driven), §11.4.58/§11.4.103 (parallel streams), §11.4.89 (background execution), and §11.4.5/§11.4.69/§11.4.107 (captured physical evidence) + §11.4 (anti-bluff). Decoupled + reusable (§11.4.28): the registry + hook + expander carry zero project-specific data and are consumed by reference (the §11.4.80 codegraph_* pattern); a consuming project ships its own registry or inherits the constitution default. Loads out-of-the-box via the §11.4.75 install seam.

Classification: universal (§11.4.17). Composes §11.4 / §11.4.5 / §11.4.6 / §11.4.17 / §11.4.20 / §11.4.28 / §11.4.35 / §11.4.58 / §11.4.66 / §11.4.69 / §11.4.70 / §11.4.75 / §11.4.80 / §11.4.89 / §11.4.91 / §11.4.93 / §11.4.103 / §11.4.105 / §11.4.107 / §11.4.109. Propagation gate CM-COVENANT-114-140-PROPAGATION (literal 11.4.140 across the consumer fleet) + recommended gate CM-ACTION-PREFIX-SYSTEM (registry exists + parses + the grammar regex matches/no-matches the canonical fixtures + the LAYER-1 mirror block is present in CLAUDE.md/AGENTS.md/QWEN.md/GEMINI.md + the BACKGROUND expansion text matches the registry verbatim + the Claude UserPromptSubmit hook + expander script exist and are executable) + paired §1.1 meta-test mutation (corrupt the BACKGROUND expansion in the registry OR strip the prefix_regex OR remove the LAYER-1 mirror block from one carrier → the gate FAILs; strip the literal 11.4.140 → the propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md §11.4.140. Non-compliance is a release blocker. No escape hatch — no --skip-action-prefix, --ignore-prefix, --no-registry, --invent-expansion-OK, --single-layer-only flag.

§11.4.141 — Token-efficiency mandate (research-derived + operator mandate, 2026-06-09)

Forensic anchor (operator mandate, 2026-06-09): the AI coding agent's token spend (input AND output) is dominated by an enormous always-loaded governance context — the consumer CLAUDE.md/AGENTS.md/QWEN.md plus the constitution submodule (measured ~170,000 tokens of static governance re-sent on EVERY turn: consumer CLAUDE.md ~90.5K + constitution/CLAUDE.md ~79.6K, a 112-row Applied-Fixes table plus verbatim §11.4.X anchor restatements) — compounded by 350–520K-token subagent transcripts, verbose status, full-file re-reads, and full-file loads where retrieval would do. The same pattern is externally documented as a scaling defect (Claude Code issue #24147 — “Cache read tokens consume 99.93% of usage quota - architectural scaling issue with CLAUDE.md re-reads”). Operator mandate: cut token spend to **30–40% of current (a 60–70% reduction)** WITHOUT degrading quality, performance, or safety, and WITHOUT breaking any existing mechanism.

Every project worked on by AI coding agents **MUST** adopt a token-efficiency regime composed of the following measures, ranked by safety and impact, and **MUST PROVE** the reduction with a measured anti-bluff harness — never an asserted number (§11.4.6 / §11.4.123). The headline 60–70% is the warm-cache target; the rule requires the *measured* best-safe reduction with cited evidence, never the estimate.

The measures (composable, each preserving every existing rule):

1. **Prompt-cache the static governance prefix (PRIMARY — the single biggest lever).** The always-loaded governance **MUST** form a byte-stable prefix with a cache breakpoint at its end, with nothing volatile rendered ahead of it. Cache reads cost $\sim 0.1 \times$ base input price (a 90% discount); writes $1.25 \times$ (5-min TTL) / $2 \times$ (1-hour). The discount applies to the dominant cost component. **Caching is transparent — the model sees byte-identical input cached or not — so it CANNOT remove a rule, weaken a gate, lower quality, or change a verdict; it changes only billing.** The only failure mode is a *silent invalidator* (timestamp/UUID/unordered JSON ahead of the breakpoint), which costs money

but never corrupts, and is detectable via `cache_read_input_tokens`. The operative discipline: **batch governance edits and keep CLAUDE.md + the constitution byte-stable mid-session** — every governance edit busts the prompt cache (BASELINE.md finding), so all governance changes in a working window are landed together rather than dribbled across turns. Claude Code caches natively; SDK callers use `cache_control:{type:"ephemeral"}`; Gemini/Qwen use their context cache via API-key auth.

2. **Subagent model-tiering + output-to-file (SECONDARY — biggest non-cache win).** Route MECHANICAL, NON-JUDGMENT subagent work (search, grep, status, doc-export, file-presence, read-only probes) to a cheaper/smaller model (Haiku-class); RESERVE the strong model for ALL reasoning, fix-design (§11.4.102), verdicts, code-review (§11.4.125), and demotion-evidence (§11.4.7). Subagents persist large output to a file and return a pointer instead of a 350–520K-token inline transcript. **The strong model owns every decision and every PASS/FAIL — the cheap model never emits a verdict — so §11.4.50 determinism and the anti-bluff covenant are untouched; quality cannot degrade.** Bounded by §12.6 (60% memory) and the §11.4.58/§11.4.103 stream caps.
3. **Thin always-loaded INDEX + on-demand detail (progressive disclosure).** Restructure the consumer governance so the always-loaded file is a concise index (one line per fix / per anchor) and the full bodies are fetched on demand. **Three invariants keep every gate green and every rule reachable: (a) each index line carries the literal 11.4.N anchor token so the CM-COVENANT-114-N-PROPAGATION literal-anchor scans still pass; (b) the canonical full text stays in the tracked, gate-scanned constitution/Constitution.md — no rule is deleted, only de-duplicated out of the consumer; (c) the full body is reachable in one hop.** This realises §11.4.35's consumer-thin / submodule-canonical split at the byte level — a de-duplication, never a deletion.
4. **Retrieval / CodeGraph-first over full-file loading.** Structural questions (where is X / what calls Y / what would break) → CodeGraph (§11.4.78/§11.4.79, already mandated + installed) or semantic search; region questions → grep-scoped read; never re-read a harness-tracked file. Read-only, already-trusted (§11.4.78) — no behaviour change.
5. **Output-token reduction.** Terse conductor status (no restatement of unchanged governance); `effort:"low"` on the mechanical-subagent allowlist only; structured output where a sink consumes it. Output is ~5× input price, so this is cost-leveraged. Presentation-only — captured-evidence artefacts (§11.4.5/§11.4.69) unchanged.
6. **Tool-call efficiency.** Batch independent tool calls in one block; reuse harness-tracked file state; reuse persisted tool outputs.
7. **Compaction / context-editing for long sessions (defensive).** Prune stale tool-results / old thinking so a 100+ turn cycle doesn't re-pay for an ever-growing transcript; never touches the governance prefix or on-disk evidence.

Mandatory measured proof (anti-bluff — §11.4.6/§11.4.50/§11.4.69/§11.4.123): the reduction MUST be proven by a token-accounting harness measuring tokens-per-development-cycle BEFORE vs AFTER on a frozen deterministic workload, split into `input_tokens` / `cache_read_input_tokens` / `cache_creation_input_tokens` / `output_tokens` from the authoritative

Anthropic usage object (NEVER tiktoken; NEVER the client-side `total_cost_usd` estimate — recompute cost from raw token fields with the published prices), reproduced N times to identical verdict (§11.4.50). Pass = AFTER $\leq 40\%$ of BEFORE (target) OR the measured best-safe reduction with a cited cold-cache reason. The AFTER run MUST ALSO show ZERO regression on the pre-build full sweep, the meta-test mutation sweep (every gate still FAILs its paired mutation — proves M3 did not turn any rule into a bluff), every propagation gate (literal anchors intact), a strong-model reasoning-path probe (proves M2/M5 tiered nothing that decides), and a cache-warm proof (`cache_read_input_tokens > 0`). Cost reduction with quality/safety regression is a §11.4 FAIL, not a win.

No measure may break or degrade any existing mechanism — and the rule is structured so none can: caching is transparent (M1); tiering and low-effort are confined to mechanical non-judgment work (M2/M5); indexing preserves the full rules by reference with literal anchors so the propagation gates pass and the canonical text stays gate-scanned (M3); retrieval uses already-trusted read-only infrastructure (M4). The §11.4 anti-bluff covenant family, §11.4.50 deterministic consistency, §11.4.125 code-review, §11.4.40 full-suite retest, and §12.6 memory ceiling are all unconditionally preserved.

Composes with §11.4.5 / §11.4.6 / §11.4.20 / §11.4.40 / §11.4.50 / §11.4.58 / §11.4.69 / §11.4.70 / §11.4.78 / §11.4.79 / §11.4.80 / §11.4.103 / §11.4.106 / §11.4.123 / §11.4.125 / §11.4.128 / §12.6 / §1.1. Classification: universal (§11.4.17) — the consuming project supplies its model IDs, its mechanical-subagent allowlist, its governance files, its structural-index tool, and its frozen workload per §11.4.35. Propagation gate CM-COVENANT-114-141-PROPAGATION (literal 11.4.141 across the consumer fleet) + recommended gate CM-TOKEN-EFFICIENCY (the measured BEFORE/AFTER harness exists + produced a verdict $\geq$ target-or-cited-floor + the non-regression sweep is green + cache-warm proof present) + paired §1.1 meta-test mutation (inject a per-turn volatile token ahead of the governance cache breakpoint → `cache_read_input_tokens` collapses → measured reduction falls below the bar → gate FAILs; restore → PASSes — proves the harness measures real caching, not a hardcoded green). Gate-code = separate work item.

Canonical authority: constitution submodule Constitution.md §11.4.141. Research + design + measurement: docs/research/token_efficiency/{RESEARCH,DESIGN,MEASUREMENT,TEST_PLAN}.md. Non-compliance is a release blocker regardless of context. No escape hatch — no `--skip-token-efficiency`, `--no-cache-governance`, `--assert-reduction-without-measuring`, `--tier-down-reasoning`, `--inline-all-governance`, `--tiktoken-estimate-OK` flag exists.

§11.4.142 — Universal code-review mandate — every change reviewed, always, no exception (User mandate, 2026-06-09)

Forensic anchor — verbatim user mandate (2026-06-09):

“ALL changes we do MUST pass through the code review step!!!
ALWAYS!!!”

EVERY change made to ANY repository this Constitution governs — without exception — MUST pass through an INDEPENDENT code-review step BEFORE it is accepted, committed, or built. There is NO change class exempt: source code, fixes, tests, gates, meta-test mutations, documentation, doc-tooling, build/CI scripts,

configuration, governance files (Constitution / CLAUDE / AGENTS / QWEN), conductor main-stream edits, sub-agent output, refactors, single-line edits — all of it. This is the ABSOLUTE form of §11.4.125 (which mandated a code-review agent gate after a batch of fixes, before the pre-build sweep + main build): §11.4.142 removes every implicit scoping that §11.4.125's "after all fixes/changes/ implementations are done" phrasing could be read to leave open — there is no "this is just a doc edit" exemption, no "this is just a one-liner" exemption, no "the author already self-reviewed (§11.4.92)" substitution, no "trivial change" carve-out. If a diff exists, it gets an independent review.

Independence is load-bearing. The review MUST be performed by a reviewer structurally separated from the author (a dedicated code-review agent, subagent-driven by default per §11.4.70 / §11.4.20, or a distinct human reviewer) — NOT the author re-reading their own work. §11.4.92's multi-pass self-evaluation is the AUTHOR-side discipline and remains required, but it does NOT satisfy §11.4.142: the whole point of an independent review seam is to catch the blind spots self-review cannot (the structurally-separated reviewer seam §11.4.70 names). A self-review labelled as the code-review step is a §11.4.142 violation of the same class as a §11.4 PASS-bluff at the review layer.

The review is itself anti-bluff. A rubber-stamp "looks good" with no examined findings is not a review (§11.4 / §11.4.1) — the review MUST genuinely analyse the change for correctness, safety (no host §12 / data §9 / target-hardware §11.4.133 regression), will-it-really-work (no solve-A-create-B per §11.4.1), end-user behaviour (§11.4 / §107), and test-genuineness (every covering test exercises the work and catches its negation per §1.1). Its conclusions are captured evidence per §11.4.5 / §11.4.69. **The review iterates to a clean GO per §11.4.134** — any finding (BLOCKING / nit / warning) re-arms the loop; the change is accepted only after a re-run review returns ZERO new findings AND ZERO warnings, with rock-solid physical evidence per §11.4.134.

Composition, not duplication. §11.4.142 is the universal "always, every change, no exception" capstone over the code-review family: §11.4.125 supplies the multi-layer review CONTENT (the six analysis dimensions + the build-gate marker CM-CODE-REVIEW-GATE-BEFORE-BUILD); §11.4.134 supplies the iterate-until-GO + rock-solid-evidence LOOP; §11.4.92 supplies the author-side self-review that PRECEDES (never replaces) the independent review; §11.4.70 / §11.4.20 supply the subagent-driven default execution model. §11.4.142 binds them and strips the last exemption: the review step is not a batch-end nicety, it is a precondition on the ACCEPTANCE of any diff. Honest boundary (§11.4.6): a passing review proves the change was independently examined and found sound to the reviewer's evidence — it does NOT replace the §11.4.108 four-layer runtime-signature verification nor the §11.4.40 full-suite retest; it is one of MULTIPLE STRONG LAYERS, and the FIRST one every change crosses.

Classification: universal (§11.4.17) — a platform-neutral acceptance-gate discipline reusable by ANY project; the consuming project supplies its concrete reviewer-dispatch mechanism + the change-acceptance seam (commit wrapper / merge queue / PR gate) per §11.4.35. Composes §11.4.1 / §11.4.4 / §11.4.5 / §11.4.6 / §11.4.20 / §11.4.40 / §11.4.69 / §11.4.70 / §11.4.92 / §11.4.108 / §11.4.110 / §11.4.125 / §11.4.134 / §107 / §1.1. Propagation gate CM-COVENANT-114-142-PROPAGATION (literal 11.4.142 across the consumer fleet) + recommended gate CM-EVERY-CHANGE-REVIEWED (every accepted change carries a fresh independent-review-completed marker for its diff, produced by a reviewer distinct from the

author, before acceptance/commit/build) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILS; accept a diff with no independent-review marker → CM-EVERY-CHANGE-REVIEWED FAILS; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md §11.4.142. Non-compliance is a release blocker regardless of context. No escape hatch — no --skip-review, --no-review, --trivial-change-exempt, --doc-edit-exempt, --self-review-suffices, --review-after-commit flag exists.

§11.4.143 — Real-user-journey mandate for video-streaming-app full-automation tests (User mandate, 2026-06-10)

Forensic anchor — verbatim user mandate (2026-06-10):

“All video streaming apps ... require to choose some title and to press proper UI button to start or resume playing! Proper UI interaction to play exact show with proper content and subtitles is MANDATORY! Without it we just eventually play on 2nd display sample, and that’s it mostly! THIS MUST BE ADDED as MANDATORY RULE regarding testing of any video streaming app in general with full automation tests! Universal in root constitution + ATMOSphere project extensions.”

Any full-automation test that asserts a video player / streaming application plays content MUST drive the REAL end-user journey through the app’s OWN user interface — launch the app → BROWSE the actual catalog → choose a SPECIFIC title → press the real Play / Resume button → confirm THAT chosen content is genuinely playing, with its correct subtitles, on the intended routing target. A test that bypasses the journey with a sample / demo / built-in-loop clip, a deep-link / am start -a VIEW / intent shortcut, a synthetic or pre-staged stream, or any path that does NOT exercise the app’s own browse-select-play UI is a §11.4 PASS-bluff at the user-journey layer: it validates ROUTING (that *something* reaches the display) while leaving the user-visible behaviour the operator actually cares about — “the show I picked actually plays” — unproven. Such shortcuts exercise neither the catalog/auth/title-resolution path nor the real Play-button → decoder → routing transition where the end-user defects live (the operator’s “we just eventually play on 2nd display sample, and that’s it mostly” is exactly this gap).

The mandate (ALL must hold): (1) **Real journey, not a shortcut** — the test drives launch → catalog browse → specific-title selection → real Play/Resume press through the app’s own UI (§11.4.48 UI-driven), NEVER a deep-link / intent / sample / loop-clip shortcut; (2) **Chosen-content confirmation** — the PASS proves the SPECIFIC selected title is playing (not merely that pixels move on the target), via the §11.4.107 liveness battery (live advancing frames + frame-advance counter + not-stale-from-previous) on the §11.4.136 real-content path, plus the §11.4.137 subtitle content-correctness oracle for the chosen title’s captions (a present cue is NOT a correct cue); (3) **Non-introspectable UIs use the pixel oracle** — when the app’s accessibility/semantic hierarchy is blank or unreliable (TV-Compose / leanback / canvas / GL streaming UIs), the test DRIVES input + ASSERTS content via the §11.4.117 CV/OCR pixel oracle (template-match the Play control + ROI-OCR the title/caption), never a hierarchy-only tool; (4) **Login via the credential single source** — apps requiring sign-in authenticate from the project’s credential single-source-of-truth (§11.4.10), never hardcoded, never logged; (5) **Honest SKIP, never a faked PASS** — where the autonomous journey is genuinely infeasible (hard human-only login / CAPTCHA, geo-block per §11.4.3, secure-surface pixel-

blanking per §11.4.112) the test is `operator_attended SKIP-with-reason` per §11.4.52 + §11.4.3 with a tracked migration item (§11.4.15 / §11.4.16) — NEVER a metadata-only / sample-played / routing-only PASS. Honest boundary (§11.4.6): “the routing fired so the title is playing” is a guess — only the chosen-content liveness + subtitle oracle on the real journey proves it; a SKIP for genuine infeasibility is correct, a PASS that never pressed the real Play button on the real title is the bluff this anchor forbids.

Classification: universal (§11.4.17) — a platform-neutral end-user-journey discipline reusable by ANY project that ships or integrates video-streaming applications; the consuming project supplies its concrete app roster, login-credential source, routing target, and UI-driving / pixel-oracle harness per §11.4.35. Composes §11.4.48 (UI-driven traversal of the real journey) / §11.4.107 (liveness — the chosen content actually advances) / §11.4.117 (CV/OCR pixel oracle for non-introspectable streaming UIs) / §11.4.136 (real-content end-to-end playback) / §11.4.137 (subtitle content-correctness for the chosen title) / §11.4.52 (autonomous-first; operator-attended only when genuinely infeasible) / §11.4.3 (topology / geo SKIP-with-reason) / §107 (end-user usability is the bar) / §1.1. Propagation gate CM-COVENANT-114-143-PROPAGATION (literal 11.4.143 across the consumer fleet) + recommended gate CM-VIDEO-REAL-JOURNEY-TEST (every video-streaming-app playback test drives the real browse-select-play journey + confirms the chosen content + subtitles, or SKIPS-with-reason) + paired §1.1 meta-test mutation (replace a real-journey test’s browse-select-play path with a deep-link / sample shortcut → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md §11.4.143. Non-compliance is a release blocker regardless of context. No escape hatch — no `--sample-playback-ok`, `--skip-real-journey`, `--deep-link-suffices`, `--routing-only-pass`, `--no-title-selection` flag exists.

§11.4.144 — Tracked/recorded-device availability-following mandate (User mandate, 2026-06-10)

Forensic anchor — direct user mandate (2026-06-10): every device the project is tracking / following / recording MUST be availability-followed with the project’s already-defined reconnection timings + the sanctioned device-recovery path; a silent recording/tracking gap — papering a disconnect as continuous capture — is a §11.4 bluff.

Every device the project is recording or tracking (test / debug / manual-testing device, across every reachable transport — USB / wireless ADB / SSH / serial / network introspection API) MUST be availability-FOLLOWED: its connection state continuously monitored, and any drop handled — never silently abandoned and never presented as a continuous recording. The always-on recorder under §11.4.128 KNOWS a tracked device is absent (its per-device loop guards on the reachable state) but, lacking a following discipline, merely spins idle — no data captured, no offline event logged, no resume, no escalation — so the recording corpus presents a continuous timeline with a silent hole. That silent hole is a §11.4 PASS-bluff at the recording-integrity layer: it claims continuous capture while no data exists for the gap.

On a tracked device leaving its reachable state the system **MUST**, automatically: (a) **DETECT** the drop and **log an honest offline event** into the recording corpus (§11.4.6 — a silent gap presented as continuous capture is a fabricated-continuity bluff; §11.4.128 — a silent recording gap defeats always-on recording); (b) **WAIT** for the device to return using the project’s ALREADY-DEFINED reconnection timings — never invented numbers (§11.4.6), the SAME grace / reconnect / poll budgets the project’s recovery path already uses; (c) **RE-ATTACH** — resume recording / tracking the moment the device returns and log an honest online / resume event; and (d) **ESCALATE** to the project’s sanctioned device-recovery path (per §11.4.69 feature class `device_recovery`) if the device does not return within the defined timeout — through the sanctioned, authorization-gated recovery entry point **ONLY**, never bypassing its gate and never performing a destructive recovery (e.g. a power-cycle) autonomously without that authorization (§11.4.21 / §11.4.101 — high-blast-radius recovery is gated; while blocked the system keeps following the device, logging the blocked-escalation honestly). A tracked-device drop that produces a silent corpus hole, a never-resumed recording, or a never-escalated permanent absence is the bluff this anchor forbids.

Classification: universal (§11.4.17) — a platform-neutral device-following discipline reusable by ANY project that records or tracks devices; the consuming project supplies its concrete tracking transport, the device set, the already-defined reconnection timings, and the sanctioned recovery entry point per §11.4.35. Composes §11.4.128 (always-on recording — §11.4.144 closes its drop-handling gap) / §11.4.69 (`device_recovery` sink-side positive evidence) / §11.4.6 (honest offline / online events, reused-not-invented timings) / §11.4.14 (watchdog children reaped on stop) / §11.4.21 + §11.4.101 (gated, non-autonomous escalation). Propagation gate CM-COVENANT-114-144-PROPAGATION (literal 11.4.144 across the consumer fleet) + recommended gate CM-DEVICE-AVAILABILITY-FOLLOWED (the recorder wires a per-device availability watchdog that reuses the defined timings + escalates only through the gated recovery path) + paired §1.1 meta-test mutation (strip the watchdog wiring / let an absence go unlogged → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule [Constitution.md](#) §11.4.144. Non-compliance is a release blocker regardless of context. No escape hatch — no --skip-availability-following, --silent-recording-gap-OK, --no-reconnect-wait, --invent-reconnect-timing, --skip-recovery-escalation, --autonomous-power-cycle-OK flag exists.

§11.4.145 — Independent multi-angle impact-research per change (User mandate, 2026-06-10)

Forensic anchor — verbatim operator intent (2026-06-10): for EVERY fix / change / new feature, independent agents **MUST** research potential issues related to that work AND connected/dependent features, from **VARIOUS ANGLES**, to be **SURE** no new issues are created, no existing working functionality is broken, and no dangerous code is introduced “like we had with the recents feature”, incorporating a **VARIETY OF TOOLS**.

Forensic anchor (FACT, project-internal): the recents / 3-button-nav path shipped a latent AIDL-interface-contract mismatch that **PASSED** code review (the diff “looked correct”) yet failed at **RUNTIME ONLY** — it ANR’d on a recents-reaping gesture — because no review angle interrogated the interface contract, the

concurrency/lifecycle, or the connected feature it silently broke; the §11.4.108 SOURCE→ARTIFACT→RUNTIME→USER-VISIBLE gap, here caught only after the fact.

§11.4.145 shifts blast-radius discovery LEFT: for EVERY change, INDEPENDENT impact-research agents (subagent-driven §11.4.70/§11.4.20, structurally separate from the author — §11.4.92 self-eval PRECEDES, never satisfies — adversarial “refute-the-change” stance to defeat the documented LLM confirmation-bias failure mode where reviewers accept plausible-but-wrong arguments) research the change AND its connected/dependent features across a CLOSED SET OF EIGHT ANGLES — (1) **correctness/logic**, (2) **regression** (what existing working feature could break — via call-graph impact enumerating every direct + transitive caller and contract-dependent feature, each mapped to its test), (3) **latent/dangerous-code — the recents class** (runtime-only failure, interface/ABI/contract mismatch, concurrency/data-race, resource lifecycle), (4) **security** (taint/injection/secret/unsafe-API), (5) **performance** (latency/memory/thermal under load vs baseline), (6) **host/data/target-hardware safety** per §12/§9/§11.4.133, (7) **cross-feature interaction** (shared state/timing/hardware contention), (8) **business-logic conformance** (matches the spec AND the connected features’ contracts, not merely compiles).

Each angle MUST name the TOOL(S) used + obtain the required DEPTH of codebase + business-logic detail (read the diff, the full changed unit, its declared contract, the spec section, and every connected feature — NEVER the diff lines alone) + emit a captured-evidence conclusion per §11.4.5/§11.4.69 (the tool’s findings/impact/diff artefact path — “looks fine” is forbidden, §11.4.1). A genuinely-N/A angle is recorded NOT_APPLICABLE: <reason> per §11.4.6, never silently skipped. The output is a per-change impact-research REPORT (one section per angle: tool + evidence path + verdict) + a single GO/NO-GO that BLOCKS acceptance/commit/build on ANY unmitigated risk in ANY angle; the change is fixed/mitigated and the affected angles re-researched, iterating to a CLEAN GO (zero unmitigated risk + zero warning, rock-solid physical evidence) per §11.4.134 before the §11.4.125 review gate.

Honest boundary (§11.4.6): a GO proves cross-angle internal-consistency + bounded blast-radius — it does NOT replace §11.4.108 four-layer runtime-signature verification nor the §11.4.40 full-suite retest; it is one of MULTIPLE STRONG LAYERS, the FIRST research pass every change crosses.

Classification: universal (§11.4.17) — the consuming project supplies its concrete toolkit (static-analyzer, call-graph engine, interface-freeze mechanism, behavioral harness) + per-angle depth per §11.4.35. Composes/strengthens §11.4.92 (adds independent tool-backed angle research) / §11.4.125 (its impact-research input) / §11.4.108 (shifts blast-radius discovery left) / §11.4.110 (angles 2/3/7 use the build-graph + interface-freeze) / §11.4.134 (iterate-to-GO) / §11.4.78 / §11.4.79 (call-graph impact engine) / §11.4.107 / §11.4.85 (behavioral angle) / §11.4.133 / §12 / §9 (angle 6) / §11.4.99 (cite latest sources) / §11.4.6 / §1.1. Propagation gate CM-COVENANT-114-145-PROPAGATION (literal 11.4.145 across the consumer fleet) + recommended gate CM-IMPACT-RESEARCH-PER-CHANGE (every accepted change carries a fresh independent impact-research report covering all eight angles with a clean GO, produced by agents distinct from the author, before acceptance/commit/build) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; accept a change with a report missing an angle or with an unmitigated NO-GO → CM-IMPACT-RESEARCH-PER-CHANGE FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md §11.4.145. Non-compliance is a release blocker. No escape hatch — no `--skip-impact-research`, `--single-angle-suffices`, `--author-researches-own-change`, `--diff-skim-OK`, `--no-go-overridable`, `--trivial-change-no-research` flag.

§11.4.146 — Reproduce-first test + same-test-confirms-fix + mandatory extend-to-all-cases workflow (User mandate, 2026-06-10)

Forensic anchor — verbatim operator intent (2026-06-10): every reported problem MUST be handled by a NAMED three-step test workflow — (1) FIRST author a test that **confirms + reproduces** the reported problem AND is **used as an investigation instrument** to obtain additional forensic data characterising the defect; (2) after the fix, use the **SAME test** to confirm the problem is GONE; (3) THEN **extend** the tests to cover ALL flows, cases, edge cases, and the variety of potential situations for that functionality, confirming no issue of any kind — with every step producing rock-solid PHYSICAL captured evidence, no false results, no bluff.

§11.4.146 does NOT re-author its component disciplines; it NAMES and BINDS them into the operator’s exact per-defect sequence and adds ONLY the two emphases the components leave implicit. The mandate (ALL hold):

(STEP 1 — REPRODUCE-FIRST + INVESTIGATE). Before any fix, author the §11.4.43 RED test as a §11.4.115 RED-baseline-on-the-broken-artifact (reproduce the defect on the CURRENT pre-fix artifact, capture defect-present physical evidence per §11.4.5/§11.4.69/§11.4.107). **NEW EMPHASIS (D1):** that same RED test is ALSO a deliberate **investigation instrument** per §11.4.102(A) — it MUST be used to gather ADDITIONAL forensic data characterising the defect (its triggers, boundaries, input/topology scope, adjacent failure modes), and that characterisation MUST feed BOTH the §11.4.102 root-cause/fix design AND the STEP-3 extend scope. A RED test used only as a binary present/absent gate, with no characterisation captured, satisfies §11.4.115 but NOT §11.4.146’s STEP 1.

(STEP 2 — SAME-TEST-CONFIRMS-FIX). After the fix, the SAME test source (the §11.4.115 polarity switch flipped `RED_MODE=1→0`) confirms the defect is ABSENT — RED-on-broken then GREEN-on-fixed, both captured, on a clean target per §11.4.108/§11.4.139, validated-first-after-redeploy per §11.4.130, deterministic per §11.4.50. No separate happy-path test substitutes for the polarity-flipped reproduction test (the §11.4.43/§11.4.115 PASS-bluff).

(STEP 3 — EXTEND-TO-ALL-CASES, mandatory per-fix). NEW EMPHASIS (D2): immediately after STEP 2 confirms the fix, the test MUST be **fanned out across the full case-space of the SAME functionality** — all flows, valid + invalid + boundary (§11.4.85 stress empty/max/off-by-one) + concurrent (§11.4.85 contention) + failure-injection (§11.4.85 chaos) + topology variants (§11.4.3) — confirming no issue of any kind, with PROVABLE enumerated coverage per §11.4.118 (a listed case-set with per-case outcome, never “no other issues found”). This is a REQUIRED step of the per-defect workflow, NOT deferred to a release-cycle discovery sweep and NOT reduced to a single guard. Each fanned-out case that exercises a user-visible path carries rock-solid physical evidence per §11.4.123;

each is registered into the §11.4.135 standing regression-guard suite so the functionality's full case-space is permanently guarded; a newly-discovered defect in the fan-out triggers §11.4.4 test-interrupt + re-enters STEP 1 for that new defect.

(ANTI-BLUFF, all steps). Every step's PASS is rock-solid CAPTURED physical evidence per §11.4.123 (§11.4.5/§11.4.69/§11.4.107) — metadata-only / config-only / absence-of-error / grep-without-runtime PASS forbidden (§11.4/§11.4.1); when STEP 1 or STEP 3 validation method is unclear, deep-research-before-declaring-untestable per §11.4.123/§11.4.8/§11.4.99. An operator-found defect the workflow's green suite missed triggers §11.4.138 (bluff-audit + permanent guard).

Honest boundary (§11.4.6): STEP 3 reduces the functionality's unknown-unknown surface; it does not prove zero remaining defects (no finite case-set can, per §11.4.118) — the earned claim is “the reproduce→confirm pair + the enumerated extend case-set are all green with captured proof,” un-exercised cases stated as honest gaps, never silently implied clean.

Classification: universal (§11.4.17) — a platform-neutral defect-lifecycle workflow reusable by ANY project; the consuming project supplies its reproduction/polarity harness, case-space enumeration, and capture mechanism per §11.4.35. Composes §11.4.43 (TDD RED step) / §11.4.115 (RED-baseline + polarity — STEP 1+2 core) / §11.4.102 (systematic-debugging — STEP 1 investigation) / §11.4.130 (validate-fix-first — STEP 2) / §11.4.108 + §11.4.139 (clean-target runtime signature) / §11.4.50 (deterministic) / §11.4.85 (stress+chaos case classes — STEP 3) / §11.4.118 (enumerated discovery coverage — STEP 3) / §11.4.135 (permanent regression guard — STEP 3 registration) / §11.4.3 (topology variants) / §11.4.123 (rock-solid-proof — all steps) / §11.4.5 / §11.4.69 / §11.4.107 (physical evidence) / §11.4.138 (operator-escape) / §11.4.4 (test-interrupt) / §107 / §1.1. Propagation gate CM-COVENANT-114-146-PROPAGATION (literal 11.4.146 across the consumer fleet) + recommended gate CM-REPRODUCE-FIRST-THEN-EXTEND (every closed defect's fix carries: a §11.4.115 reproduce-first polarity test with captured defect-characterisation [STEP 1], its RED_MODE=0 GREEN confirmation [STEP 2], AND an enumerated per-functionality extend case-set registered into the §11.4.135 suite [STEP 3] — a closure with the reproduce→confirm pair but NO enumerated extend case-set FAILs) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; close a defect with only the reproduce→confirm pair and no extend-case-set → CM-REPRODUCE-FIRST-THEN-EXTEND FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md §11.4.146. Non-compliance is a release blocker. No escape hatch — no --skip-reproduce-first, --fix-without-red, --skip-extend-to-all-cases, --reproduce-confirm-suffices, --defer-extend-to-release-sweep, --single-guard-suffices, --no-defect-characterisation flag.

§11.4.147 — Crashed-agent respawn-until-complete + no-work-loss registry mandate (User mandate, 2026-06-10)

Forensic anchor — verbatim operator intent (2026-06-10): “any agent that crashed because of something MUST BE respawned and finish its work at some point! We MUST NOT lose any work, forget about or have it corrupted!” Forensic case study (FACT, this session): dispatched subagents died on **transient** causes — API Error: Server is temporarily limiting requests (rate-limit) killed 5 at once, API Error: socket connection closed unexpectedly killed 2, and

one left PARTIAL edits (two source files written, the dependent SystemUI sliders + doc + test NOT done). Each was re-dispatched by pure conductor vigilance with **no mechanical guarantee** — the moment conductor context is lost / compacted / the conductor itself crashes, an in-flight-but-dead work unit is silently dropped: work lost / forgotten / corrupted.

Every agent/subagent dispatched to perform a work unit **MUST** be tracked through its full lifecycle so that a crash **NEVER** loses, forgets, or corrupts its work. A crashed agent is **NOT** a completed agent (§11.4.6 — “it probably finished before dying” is a forbidden guess; an abnormal termination is positive evidence the work unit is OPEN). The mandate (ALL hold):

(a) Mandatory work-unit / agent REGISTRY. Every dispatched agent/subagent **MUST** have a durable, append-only, machine-readable registry entry carrying at minimum: a stable agent/work-unit id, its task + file-scope (§11.4.58 PWU manifest), its declared output path(s) + captured-evidence dir (§11.4.5/§11.4.69), its dispatch timestamp, and a status from the CLOSED SET {dispatched | in-flight | crashed | respawned | complete}. The registry reuses the §11.4.116 sync substrate (append-only JSONL event stream + atomically-rewritten status snapshot) — agent-lifecycle transitions are events; the §11.4.116 integrity property holds (“an entry with no dispatch-event cannot show complete”). The registry is the SINGLE SOURCE OF TRUTH for “is this work owed?” — an unregistered dispatch is itself a §11.4.147 violation (an untracked agent cannot be respawned).

(b) Mechanical CRASH-DETECTION + RESPAWN-until-complete. Any abnormal agent termination — transient API/rate-limit error (Server is temporarily limiting requests), socket-close (socket connection closed unexpectedly), any non-completion exit, OR a terminal error after the runtime’s own retries are exhausted — **MUST** flip the registry entry to crashed and keep the work unit OPEN. An OPEN unit **MUST** be respawned (a fresh agent claims the same id + scope + preserved state) and re-respawned until an agent reaches complete. Respawn is the safe, reversible, bounded-blast-radius decision the conductor takes autonomously per §11.4.101 (never blocks the loop waiting for a human). Transient causes use the project’s ALREADY-DEFINED backoff/retry budgets — never invented numbers (§11.4.6); a genuinely-non-transient terminal error (deterministically reproducible after respawn) is investigated per §11.4.102 before further respawn, never spun in a blind retry loop.

(c) PARTIAL-STATE handling — preserve, then resume-or-clean-restart. A crashed agent’s uncommitted edits + output doc **MUST** be PRESERVED (never silently discarded → work lost; never blindly committed → corruption). Before resuming, the respawn **MUST** run the §11.4.84 quiescence check on the preserved tree (every modified file accounted-for against the unit’s scope; no mutation/// always pass/_mutated_* residue; no half-written/torn artifact). Then **EITHER** (i) **RESUME idempotently** from the preserved partial when it passes the check (consistent + complete-enough to continue), OR (ii) **CLEAN-RESTART** the unit from a known-good base (§9.2 pre-op backup, reversible per §11.4.101) when the partial is inconsistent / incomplete-unsafely / fails the §11.4.84 check. Either path guarantees nothing is lost and nothing is corrupted. Per-agent git worktree isolation (§11.4.58 L4 / §11.4.84) keeps a crashed unit’s partial tree from contaminating other streams.

(d) COMPLETION criterion. A work unit is **DONE only** when an agent reaches complete with its required captured evidence/output landed (§11.4.5/§11.4.69 + the §11.4.116 verdict-carries-evidence-path rule). The endless-loop done-condition (§11.4.87 / §11.4.94 / §11.4.97 / §11.4.126) **MUST NOT** read as satisfied while any registry entry is dispatched / in-flight / crashed / respawned-not-yet-complete; the zero-idle survey (§11.4.94) **MUST** treat every non-complete entry as an OPEN actionable item to reclaim. A registry showing complete without the landed evidence is a §11.4 PASS-bluff at the agent-lifecycle layer.

Honest boundary (§11.4.6): the registry + respawn guarantee that work is *not lost or corrupted*; it does NOT itself prove the work is *correct* — the respawned unit's output still crosses §11.4.108 four-layer verification, the §11.4.125/§11.4.142 code-review gate, and §11.4.40 full-suite retest. §11.4.147 is the durability layer beneath those, not a substitute. §11.4.147 is the agent-side analogue of §11.4.144 (device availability-following): same detect-drop → wait/backoff → re-attach/respawn → escalate-if-non-transient shape — devices for §11.4.144, dispatched agents for §11.4.147.

Classification: universal (§11.4.17) — the consuming project supplies its concrete registry path/format, its crash-detection signal source (runtime exit codes / sync-stream error events), its respawn-dispatch mechanism, and its already-defined backoff budgets per §11.4.35. Composes §11.4.6 (no-guessing — crash ≠ done) / §11.4.58 (PWU = the work unit + worktree isolation) / §11.4.84 (partial-state quiescence check before resume) / §11.4.87 + §11.4.94 + §11.4.97 + §11.4.126 (loop done-condition + zero-idle reclaim) / §11.4.101 (autonomous respawn-vs-restart decision, reversible-with-backup) / §11.4.116 (registry reuses the sync-channel substrate) / §11.4.102 (non-transient terminal error → systematic-debugging before blind respawn) / §11.4.108 + §11.4.125 + §11.4.142 + §11.4.40 (the respawned output still crosses every downstream verification layer) / §9.2 (clean-restart pre-op backup) / §11.4.128 + §11.4.144 (the recording/device-following siblings — same detect→wait→re-attach→escalate shape, agent-side here). Propagation gate CM-COVENANT-114-147-PROPAGATION (literal 11.4.147 across the consumer fleet) + recommended gate CM-CRASHED-AGENT-RESPAWN-TRACKED (every dispatched agent has a registry entry; every crashed entry is respawned-until-complete; the loop done-condition reads no non-complete entries; partial state preserved + §11.4.84-checked before resume) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; mark a crashed entry as the loop's done-condition complete without landed evidence, OR drop a dispatched agent without a registry entry → CM-CRASHED-AGENT-RESPAWN-TRACKED FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md §11.4.147. Non-compliance is a release blocker. No escape hatch — no --skip-agent-registry, --crash-equals-done, --no-respawn, --discard-partial-state, --blind-commit-partial, --forget-dead-agent, --loop-done-with-crashed-entries flag.

§11.4.148 — Workable-item integrity (status+type+id) + comprehensive structured description + bidirectional external-tracker sync + BLOCKED unblock-choices mandate (User mandate, 2026-06-10)

Forensic anchor — direct operator mandate (2026-06-10): every workable item MUST carry a VALID status, a VALID type, and a stable unique identifier at all times (never one without all three); every item MUST carry a comprehensive, structured, self-explanatory description (what it is, how it manifests, how to reproduce, acceptance criteria); a BLOCKED item MUST state WHY it is blocked AND the choices that would unblock it; the single source of truth (the SQLite workable-items DB), the rendered docs, and the external tracker (e.g. ClickUp) MUST be regularly, never-missed, fully synchronized in BOTH directions through the mechanical docs-sync engine; the external-tracker push MUST carry statuses, types, assignee, and sub-tasks, default an unset assignee from an env var, and be idempotent.

This anchor BINDS + STRENGTHENS the existing workable-item discipline (§11.4.15 status / §11.4.16 type / §11.4.54 ATM-NNN id / §11.4.21 Operator-blocked / §11.4.91 description-clarity / §11.4.93 + §11.4.95 SQLite SSoT / §11.4.106 docs_chain) into ONE integrity contract spanning the three surfaces — DB [?] docs [?] external tracker — and adds the operator's three new emphases (D1 the no-item-without-all-three invariant, D2 the comprehensive structured description, D3 the BLOCKED unblock-CHOICES tightening), citing rather than re-authoring its components. The mandate (ALL hold):

(D1) No item without a valid status + valid type + stable id — across ALL three surfaces. Every workable item, at every moment, MUST carry (i) a ****Status:**** from the §11.4.15 / §11.4.21 / §11.4.90 closed set, (ii) a ****Type:**** from the §11.4.16 closed set {Bug | Feature | Task}, and (iii) a stable, unique, monotonic, append-only [<ID-NNN>] identifier per §11.4.54. An item missing ANY of the three in the DB, in the rendered docs, OR in the external tracker is a §11.4.148 violation — the validator FAILs (non-zero exit, release-blocker). The id is the binding key that keeps the same item identifiable across DB rows, doc headings, and tracker tasks; an id present on one surface but absent on another is a sync-integrity failure.

(D2) Comprehensive structured description per item. Every item MUST carry a comprehensive, self-contained, structured description — at minimum WHAT it is (the §11.4.91 ≥6-word / ≥40-char clear meaning), HOW it manifests (observed user-visible / system behaviour), HOW to reproduce (deterministic steps where applicable, §11.4.115 / §11.4.146 reproduce-first), and ACCEPTANCE CRITERIA (the captured-evidence verdict that closes it per §11.4.123 / §11.4.69). The description lives in the §11.4.93 DB description column AND renders into the docs AND pushes to the external tracker — never a one-liner stub, never a §11.4.91 anti-pattern fragment. A “fix-it-later”/placeholder description is a §11.4 PASS-bluff at the planning layer (the operator / a future agent / a teammate cannot understand the item).

(D3) BLOCKED items carry WHY + the unblock CHOICES. Tightens §11.4.21: every item whose status is Operator-blocked (incl. any Blocked/BLOCKED input alias normalised to the canonical Operator-blocked value — documented alias, never a silent fork, §11.4.6) MUST carry, beyond the §11.4.21 Operator-Block-Details WHAT/WHY/WHO, an enumerated set of UNBLOCK CHOICES — the

closed list of decisions/actions that would unblock it
(e.g. [A] grant access X · [B] approve removal of Y · [C] provide credential Z, mirroring the §11.4.66 2–4-option interactive-clarification shape). The validator FAILs a BLOCKED item whose unblock field is empty OR carries no enumerated choice marker — “blocked, no way forward stated” is forbidden, the agent / operator must always know the choices that clear the block.

(D4) Bidirectional DB [?] docs [?] external-tracker sync, never missed, docs_chain-bound. The §11.4.93 / §11.4.95 git-tracked SQLite DB is the SINGLE SOURCE OF TRUTH; the rendered docs (Issues / Issues_Summary / Fixed / Fixed_Summary / per-item docs + their §11.4.65 / §11.4.106 exports) and the external tracker are its DERIVED surfaces. A full synchronization pass MUST run regularly and on every state change so no surface drifts: DB→docs (db→to→md), docs→DB (md→to→db, both byte-identical round-trip per §11.4.93), DB→external-tracker push, with a drift-proof §11.4.86 fingerprint (sha256 of the sorted item keyset, NOT mtime) gating freshness so a forgotten sync FAILs a pre-build gate. The sync is bound into the §11.4.106 docs_chain engine as a registered consumer context (never an ad-hoc parallel script) so drift is mechanically caught, never vigilance-dependent.

(D5) Generic external-tracker sync — statuses/types/assignee/sub-tasks, default-assignee env var, idempotent. The external-tracker push MUST carry: each item’s status (collapsed onto the tracker’s native status set per §11.4.33 when the tracker’s status vocabulary is fixed — the API-cannot-create-statuses case is a §11.4.112 structural constraint, the precise value preserved in a description header, never lost), its type, its assignee (a §11.4.104 participant handle — created_by / assigned_to), and its sub-tasks. An UNSET assignee defaults from a project-supplied env var (the §11.4.104 operator-username pattern), never hardcoded, never logged (§11.4.10). The push MUST be IDEMPOTENT — a dry-run-then-real, match-by-stable-key ([<ID>] name prefix / id custom-field) plan, present⇒UPDATE / absent⇒CREATE, rate-limited, credential-redacted, with a captured sink-side created=N updated=M failed=0 proof per §11.4.69. The tracker-sync MACHINERY stays project-agnostic (§11.4.28) — the consumer registers its concrete tracker (service, list/board id, field map) at runtime via the public API; the constitution names the discipline, not the vendor.

Anti-bluff (§11.4): every sync run + every validator pass carries captured evidence per §11.4.5 / §11.4.69 (the round-trip diff, the fingerprint match, the tracker created/updated/failed line under qa-results/); a sync claimed-run with no growing evidence trail, a complete/Fixed item with no acceptance-criteria description, or a BLOCKED item with no unblock choices are all PASS-bluffs at the workable-item-integrity layer. Honest boundary (§11.4.6): the integrity contract guarantees every item is well-formed + the three surfaces agree, NOT that the item’s underlying work is correct — that still crosses §11.4.108 / §11.4.40 / §11.4.123.

Classification: universal (§11.4.17) — the consuming project supplies its concrete DB path, id prefix, external-tracker service + list/board id + field map + default-assignee env var, and docs_chain context per §11.4.35. Composes §11.4.15 / §11.4.16 / §11.4.21 / §11.4.33 / §11.4.34 / §11.4.54 / §11.4.66 / §11.4.86 / §11.4.91 / §11.4.93 / §11.4.95 / §11.4.104 / §11.4.106 / §11.4.112 / §11.4.123 / §11.4.10 / §11.4.28 / §11.4.69 / §11.4.6 / §1.1. Propagation gate CM-COVENANT-114-148-PROPAGATION (literal 11.4.148) + recommended gates CM-ITEM-INTEGRITY-STATUS-TYPE-ID (every item carries valid status+type+id on every surface) / CM-ITEM-COMPREHENSIVE-DESCRIPTION (every item’s description has the WHAT/manifest/repro/acceptance structure) / CM-BLOCKED-UNBLOCK-CHOICES (every

BLOCKED item enumerates unblock choices) / CM-TRACKER-SYNC-IDEMPOTENT (the external-tracker push is bidirectional-fingerprinted + idempotent + docs_chain-bound) + paired §1.1 meta-test mutations (strip the literal → propagation gate FAILs; drop an item's type / id / description-structure / unblock-choices, OR break the tracker-sync idempotency key → the respective gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md §11.4.148. Non-compliance is a release blocker. No escape hatch — no --item-without-status, --item-without-type, --item-without-id, --stub-description-OK, --blocked-without-choices, --skip-tracker-sync, --one-way-sync-OK, --non-idempotent-tracker-push, --hardcode-assignee flag.

§11.4.149 — Per-workable-item testing-diary mandate (User mandate, 2026-06-10)

Forensic anchor — direct operator mandate (2026-06-10): every workable item MUST carry a TESTING DIARY — a chronological, append-only record of every test run against it, each entry capturing date/time, who tested (tested-by), the result, in-depth observations (the technical background a future fix needs), and the action taken + why (whether the run changed the item's status, and the reasoning); both an in-depth diary AND an at-a-glance summary; modelled on the external tracker as a sub-task of the item with a {TODO | In-progress | Completed} lifecycle; persisted in the SQLite single source of truth + exported to all required formats + pushed to the external tracker; bound into the mechanical docs-sync engine so it is never missed; minimal-LLM (deterministic tooling, the prose authored by whoever ran the test); under 100% test-coverage discipline.

The testing diary is the evidentiary HISTORY layer for a workable item — distinct from the §11.4.93 item_history (which records lifecycle STATE transitions) and from the §11.4.55 Reopens.md (which aggregates reopen cycles): the diary records TEST EXECUTIONS, each of which may or may not cause a state transition. The mandate (ALL hold):

(a) Append-only per-item diary table in the SQLite SSoT. A test_diary table (additive to the §11.4.93 / §11.4.95 schema, never a competing source) records one row per test run, soft-keyed to the item's stable id, with at minimum: ISO-8601-UTC date_time (sortable, deterministic §11.4.50); tested_by from a closed set (e.g. User | Operator | AI-agent | HelixQA); result from the §11.4.45 closed verdict vocabulary {PASS | FAIL | SKIP} + a short detail; in-depth observations (long markdown — facts or explicit UNCONFIRMED:/ PENDING_FORENSICS: per §11.4.6, never “likely”; this is the background a future fix relies on); action_taken + a status_changed flag + from/to (did this run change the item's status, and WHY/why-not); and a §11.4.69 evidence_path + feature_class — with a structural constraint that a PASS row WITHOUT a non-empty evidence path is impossible (a PASS-bluff is rejected by the schema itself).

(b) In-depth diary + at-a-glance summary. Per item, BOTH an in-depth diary document (one section per run, full observations + action + evidence link) AND a derived at-a-glance summary (a rollup VIEW — total/pass/fail/skip runs, last verdict, last run, status-change count, distinct testers, distinct feature-classes — DERIVED, never a duplicate source per §11.4.93). The summary feeds the §11.4.132 risk-ordering inputs (fail-runs / status-changes / last-result) without re-reading every observations blob.

(c) Four-format per-item exports, docs_chain-bound, never missed. Each item's Diary + Diary_Summary render to the project's full export-format set (§11.4.65 + any project-added format) with the freshness invariant (derived-format mtime ≥ source mtime), wired through the §11.4.106 docs_chain engine as a registered consumer context with a §11.4.86 drift-proof fingerprint (sha256 of the sorted diary keyset, NOT mtime) so a diary row added but not exported/pushed FAILs a pre-build gate — the never-missed mechanism.

(d) External-tracker sub-task model with lifecycle. On the external tracker, the workable-item task keeps its description = the item (§11.4.148 D2 — NOT polluted with diary text); each diary run is a CHILD SUB-TASK of that task (type task), carrying a {TODO | In-progress | Completed} lifecycle status (collapsed onto the tracker's native status set per §11.4.33 / §11.4.112 when the tracker's vocabulary is fixed), its observations + action + an Evidence: line (the path, not the raw artefact — §11.4.10 / §11.4.13), and a stable diary-entry idempotency key. The push reuses the §11.4.148 D5 idempotent, rate-limited, credential-redacted, sink-side-proven discipline; the sub-task mapper stays project-agnostic (§11.4.28).

(e) Minimal-LLM, anti-bluff, 100% test-covered. The diary tooling (add / export / tracker-sync / validate) is deterministic bash/Go with ZERO LLM in the data path — the observations prose is authored by whoever ran the test; the tooling only stores/renders/pushes/validates it. The diary is covered per §11.4.27 by every applicable test type (unit + integration-against-the-real-tracker with an honest §11.4.3 SKIP-when-token-absent, never a faked PASS + export + HelixQA Challenge + paired §1.1 mutation) so a diary PASS-bluff is mechanically impossible. Honest boundary (§11.4.6): the diary guarantees a complete, auditable test-history per item, NOT that any single run's verdict is correct — that rests on the run's own captured evidence per §11.4.69 / §11.4.107 / §11.4.123.

Classification: universal (§11.4.17) — the consuming project supplies its concrete DB path, diary doc layout, export formats, external-tracker sub-task field map, and docs_chain context per §11.4.35. Composes §11.4.6 / §11.4.27 / §11.4.45 / §11.4.50 / §11.4.55 / §11.4.65 / §11.4.69 / §11.4.86 / §11.4.93 / §11.4.95 / §11.4.106 / §11.4.107 / §11.4.123 / §11.4.132 / §11.4.148 / §11.4.10 / §11.4.13 / §11.4.28 / §11.4.33 / §11.4.112 / §1.1. Propagation gate CM-COVENANT-114-149-PROPAGATION (literal 11.4.149) + recommended gates CM-TEST-DIARY-SYNC (the diary fingerprint gate — a diary row not exported/pushed FAILs) / CM-DIARY-PASS-REQUIRES-EVIDENCE (a PASS diary row carries a non-empty evidence path) + paired §1.1 meta-test mutations (strip the literal → propagation gate FAILs; strip the evidence-required constraint OR the sub-task lifecycle / fingerprint → the respective gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md §11.4.149. Non-compliance is a release blocker. No escape hatch — no --skip-testing-diary, --diary-without-evidence, --no-diary-summary, --diary-without-tracker-subtask, --llm-in-diary-data-path, --diary-export-optional flag.

§11.4.150 — Mandatory deep multi-angle web research per change/issue, before declaring fixed or structural (User mandate, 2026-06-11)

Forensic anchor — verbatim user mandate (2026-06-11):

“For every single issue we fix or improvement we make — besides tight systematic-debugging, fixing, code review by independent agents, comprehensive tests — we MUST ALWAYS do deep web research from various angles! No matter how big/small/simple/complex, dig deep the internet for articles, technical documentation, APIs, open-source code — EVERYTHING that can help make the best possible solution OR confirm we don’t have a more serious problem we’re unaware of! We MUST do everything possible to STOP the constant issue-reopening and finally start closing items as fixed+working! Do this ALWAYS in parallel with the main work stream! Add the mandatory rules into the constitution Submodule CONSTITUTION.md, CLAUDE.md, AGENTS.md, QWEN.md and other relevant files. Commit + push to all upstreams. Apply ASAP to every workable item.”

Forensic case study (FACT, 2026-06-11). A subtitle-on-secondary-display goal had been verdicted Won't-fix: structurally-impossible (the §11.4.112 secure-surface pixel-blanking reasoning). A subsequent DEEP multi-angle web research pass OVERTURNED that verdict — it surfaced the real, workable OCR-of-the-PRIMARY-display path (read the caption on the introspectable primary surface rather than the FLAG_SECURE secondary), proving the goal was NOT structurally impossible, only mis-classified for want of research. A single research pass from a different angle converted a “we can’t” into a shipped capability. This is the exact reopen-cycle the mandate exists to break: items churn between Fixed[?]Reopened, or sit mis-labelled structural/won't-fix, because the solution-space (or the confirmation that the closure is genuinely safe) was never deep-researched from enough angles.

For EVERY fix / improvement / change / closure — no matter how big, small, simple, or complex — the agent MUST, IN ADDITION to tight systematic-debugging (§11.4.102), the fix itself, independent-agent code review (§11.4.125 / §11.4.142), and comprehensive multi-layer tests (§11.4.4(b) / §11.4.40), perform a DOCUMENTED **deep multi-angle web research pass** that digs the internet from VARIOUS angles — articles, official + vendor technical documentation, API references, standards, issue trackers, maintainer guidance, reusable open-source code — to BOTH (i) discover the best possible solution AND (ii) confirm there is NOT a more serious underlying problem the team is unaware of. This BINDS + STRENGTHENS the existing research/proof anchors into one mandatory, reopen-breaking, run-in-parallel discipline; it does not re-author them:

- **(A) No closure-as-fixed/structural WITHOUT a documented deep-research pass.** No issue may be marked Fixed/Implemented/Completed (§11.4.33) OR classified structurally-impossible won't-fix (§11.4.112) UNTIL a deep multi-angle research pass is documented for it. §11.4.8 already requires deep research before designing a non-trivial fix and §11.4.123 already requires it when validation is unclear; §11.4.150 makes the pass UNCONDITIONAL — required for EVERY workable item, however trivial, at the closure AND structural-verdict gates — precisely because the forensic case proved a “trivial-looking” structural verdict was wrong.
- **(B) Multiple angles, not a single lookup.** The pass MUST interrogate the problem from several distinct angles (e.g. official-docs / standards / known-bug-trackers / alternative-approach / failure-mode / security / performance / platform-constraint / open-source-precedent — the consuming project picks the applicable subset, ≥ 2 genuinely-distinct angles), so a single-source confirmation-bias miss (the §11.4.145 LLM-confirmation-bias failure mode) cannot pass as research. A one-link drive-by is NOT a deep multi-angle pass.

- **(C) Confirm-no-bigger-problem, not just find-a-fix.** The pass **MUST** explicitly seek evidence that the chosen fix does NOT mask a deeper defect AND that the closure (or the structural verdict) is genuinely safe — the “do we actually have a more serious problem we’re unaware of?” check. Absence of a found bigger problem is recorded as a positive finding only when the search that would have surfaced it is enumerated (composes §11.4.118 discovery-completeness — “we found nothing worse” requires “here is what we searched”).
- **(D) Latest-source + cited.** Sources **MUST** be the **LATEST** authoritative versions per §11.4.99 (never training-data / memory), each cited by URL + access date in the item’s research artefact AND the closure commit footer (Deep-research <date>: <urls> OR the literal NO external solution found – original work per §11.4.8 when the search genuinely yields nothing). Stale or memory-sourced “research” is a §11.4.99 + §11.4.150 violation.
- **(E) Reopen-breaking is the PURPOSE.** The explicit goal is to STOP the constant Fixed☐Reopened churn and finally close items as fixed-AND-working: a reopen (§11.4.34 / §11.4.55) whose root cause a deep multi-angle pass would have surfaced is a §11.4.150 miss; and a structural/won't-fix verdict reached without the pass is forbidden (the case study’s overturned verdict is the canonical anti-pattern). Composes §11.4.7 (demotion-evidence — a reopen IS a demotion, and the research is part of the demotion-prevention) + §11.4.112 (structural-impossibility now **REQUIRES** the deep-research pass + its cited authorities before the verdict is earned).
- **(F) ALWAYS in parallel with the main stream.** The research pass **MUST** run as a background, subagent-driven work stream (per §11.4.70 / §11.4.20 / §11.4.103 continuous-parallel-stream + §11.4.89 background execution) concurrent with the main fix/build/test work — it **NEVER** serialises the main stream or stalls the loop (§11.4.94 / §11.4.97 / §11.4.126). A blocked research angle parks that angle, it does not pause the closure pipeline (§11.4.101).
- **(G) Apply ASAP to every workable item.** The discipline applies retroactively + going forward to EVERY workable item (§11.4.93 / §11.4.95 SSoT); items closed without a documented pass are re-audited as part of the §11.4.40 / §11.4.42 release-gate sweep.

Honest boundary (§11.4.6): a deep-research pass reduces the unknown-unknown surface and breaks the most common reopen causes — it does NOT prove zero remaining defects (§11.4.118) and does NOT replace §11.4.108 four-layer runtime-signature verification, §11.4.125 / §11.4.142 independent review, or §11.4.40 full-suite retest. It is one of MULTIPLE STRONG LAYERS — the research layer that every fix/closure/structural-verdict additionally crosses. “We probably don’t have a bigger problem” without the enumerated multi-angle search is a guess (§11.4.6), never a finding.

Classification: universal (§11.4.17) — a platform-neutral research-and-closure discipline reusable by ANY project; the consuming project supplies its concrete research corpora, angle set, item tracker, and closure-commit-footer convention per §11.4.35. Composes §11.4.8 (deep-web-research-before-implementation — §11.4.150 makes it unconditional + multi-angle + closure-gating) / §11.4.99 (latest-source verification of the cited authorities) / §11.4.123 (rock-solid-proof-or-deep-research — §11.4.150 raises “when unsure” to “always”) / §11.4.118 (discovery-completeness — the confirm-no-bigger-problem search is enumerated) / §11.4.145 (independent multi-angle impact-research — §11.4.150 is its external-research sibling, adversarial against confirmation bias) / §11.4.125 / §11.4.142 (independent review — research precedes, never substitutes) / §11.4.7 (demotion/reopen

evidence) / §11.4.112 (structural-impossibility now requires the cited-authorities pass) / §11.4.34 / §11.4.55 (reopen attribution — the churn this breaks) / §11.4.70 / §11.4.20 / §11.4.89 / §11.4.103 (parallel background subagent stream) / §11.4.40 / §11.4.42 (release-gate re-audit) / §11.4.93 / §11.4.95 (per-item SSoT). Propagation gate CM-COVENANT-114-150-PROPAGATION (literal 11.4.150) + recommended gate CM-DEEP-RESEARCH-PER-ISSUE (every closed/structural-verdicted item carries a documented multi-angle deep-research artefact + cited-source closure footer, run in parallel, before the closure/structural verdict is accepted) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; close an item or reach a structural verdict with no documented multi-angle research artefact / cited footer → CM-DEEP-RESEARCH-PER-ISSUE FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md §11.4.150. Non-compliance is a release blocker. No escape hatch — no `--skip-deep-research`, `--single-source-suffices`, `--trivial-change-no-research`, `--close-without-research`, `--structural-without-research`, `--serialise-research`, `--research-from-memory-OK` flag.

§11.4.151 — Project-prefixed release-tag/version-naming mandate (User mandate, 2026-06-12)

Forensic anchor — verbatim user mandate (2026-06-12):

“Every release tag and version name we create — on the main repository and on every Submodule we own — MUST be prefixed with the project’s release prefix, e.g. `helix_translate-1.0.0-dev-0.0.1`. The prefix MUST come from `HELIX_RELEASE_PREFIX` in our `.env` if it is set, otherwise from the lowercased project root directory name. The SAME prefix MUST be used across the main repo and all owned Submodules in one release so a release is greppable across every repository.”

Every release tag AND every version name created on the main repository AND on every owned-by-us submodule (per §11.4.28) MUST be prefixed with the project’s release prefix, in the form `<PREFIX>-<version>` (canonical example: `helix_translate-1.0.0-dev-0.0.1`). The prefix makes a release identifiable + greppable across every repository the release spans, so a single `git tag -l '<PREFIX>-*'` (or equivalent) enumerates the whole release surface in one query.

Prefix resolution order (closed-set, deterministic — §11.4.6 no-guessing):

1. **HELIX_RELEASE_PREFIX from the project’s .env** — authoritative when set. `.env` is git-ignored per §11.4.30; the variable MUST be documented in the tracked `.env.example` placeholder so a fresh clone knows the contract (the regeneration-mechanism discipline of §11.4.77 applies — `.env` is re-obtained, never committed).
2. **Fallback = the lowercased snake_case form of the project root directory name** (no spaces), per §11.4.29 — used whenever `HELIX_RELEASE_PREFIX` is unset/empty. The fallback is deterministic from the checkout, so a release prefix is always resolvable without operator input.

The prefix MUST be IDENTICAL across the main repo and all owned submodules within a single release — a release that tags the main repo `<PREFIX>-1.2.0` while tagging an owned submodule `<OTHER>-1.2.0` (or unprefixed) is a §11.4.151 violation: the cross-repo `grep` no longer enumerates the

release. Version codes increment monotonically within the prefix (<PREFIX>-...-0.0.1 → <PREFIX>-...-0.0.2 → ...), never reset, never skipped — same append-only discipline as §11.4.54's ticket ids applied to the version axis.

Honest boundary (§11.4.6): the prefix guarantees a release is identifiable + uniform across every repository, NOT that the release contents are correct — the tag is still only created after the §11.4.40 full-suite retest GREEN and reaches every upstream via the §11.4.113 merge-onto-latest-main path (NEVER a force-push), fanned out per §2.1.

Classification: universal (§11.4.17) — a platform-neutral release-naming discipline reusable by ANY project; the consuming project supplies its concrete prefix value + the HELIX_RELEASE_PREFIX env var per §11.4.35. Composes §2.1 (multi-upstream push — the prefixed tag fans out to every remote) / §11.4.29 (lowercase snake_case — the fallback prefix follows it) / §11.4.30 (.env git-ignored — the authoritative prefix source) / §11.4.40 (full-suite retest before the tag is created) / §11.4.113 (absolute no-force-push — the prefixed tag reaches upstreams via merge-onto-latest-main) / §11.4.126 (the release-scope terminal condition is a published, prefixed tag). Propagation gate CM-COVENANT-114-151-PROPAGATION (literal 11.4.151 across the consumer fleet) + recommended gate CM-RELEASE-PREFIX-NAMING (every release tag/version on the main repo + every owned submodule carries the resolved <PREFIX>- prefix, identical across the release) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; create an unprefix or differing-prefix release tag → CM-RELEASE-PREFIX-NAMING FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md §11.4.151. Non-compliance is a release blocker. No escape hatch — no --no-release-prefix, --unprefixed-tag, --prefix-optional, --differing-submodule-prefix flag.

§11.4.152 — Crashlytics-recorded-data continuous monitoring + systematic-debug + regression-test-coverage mandate (User mandate, 2026-06-13)

Forensic anchor — verbatim user mandate (2026-06-13):

“For every project that has Firebase Crashlytics enabled / wired, we MUST continuously monitor ALL of the Crashlytics-recorded data — crashes, ANRs, performance traces, and non-fatals — systematically debug each, fix and improve, and cover everything with validation and verification tests. This MUST be checked regularly, with no false results and no bluff of any kind!”

Every project that has Firebase Crashlytics enabled / wired (the SDK linked into a shipping artifact, crash + non-fatal + ANR reporting active) MUST treat the Crashlytics console as a first-class captured-evidence channel from real end-user devices and continuously process every datum it records. Crashlytics IS the user-visible-behaviour signal §11.4 exists to protect: a recorded crash / ANR / performance regression / non-fatal is a real end user hitting a broken feature, and a recorded issue left unmonitored, undebugged, unfixed, or uncovered-by-test is the precise “tests pass while the feature is broken in the wild” failure mode the §11.4 covenant forbids — an open Crashlytics issue with a green test suite is a §11.4 PASS-bluff at the field-telemetry layer.

§11.4.152 STRENGTHENS + COMPLETES §11.4.47 (Firebase data review mandate): §11.4.47 owns the periodic REVIEW + dedup + Issue-creation pass that surfaces Crashlytics/Analytics/Performance findings and maps them to tracker entries; §11.4.152 owns what happens to each surfaced item AFTER it is recorded — the systematic-debug → fix/improve → regression-test-coverage lifecycle that drives it to a proven, regression-immune closure. §11.4.47 finds it; §11.4.152 fixes it and proves it stays fixed. Both are mandatory; neither substitutes for the other.

The mandate (ALL must hold):

1. **Continuous monitoring of ALL four Crashlytics surfaces.** The monitoring pass MUST cover fatal crashes, ANRs, performance traces / regressions, AND non-fatals — all four. Skipping any one surface is a §11.4 PASS-bluff (a regression surfaces on the skipped axis and the team never sees it — the §11.4.47 three-source lesson extended to the non-fatal + ANR + performance axes). Non-fatals in particular are the silent class: every catch / fallback / recovered-error path that records a non-fatal is a real degraded user experience and MUST be triaged, not ignored because the app did not crash.
2. **Systematic-debugging of each recorded issue (reproduce-before-fix).** Every Crashlytics-recorded issue selected for fix MUST go through the §11.4.102 systematic-debugging arc — Iron Law: NO FIX WITHOUT ROOT-CAUSE INVESTIGATION FIRST — root-cause → pattern → falsifiable hypothesis proven with captured evidence → fix designed against the proven cause. The §11.4.115 RED-baseline-on-the-broken-artifact reproduces the recorded defect (the recorded stacktrace / ANR trace / non-fatal context IS the §11.4.5/§11.4.69 captured evidence of defect-present) BEFORE the fix; a fix without a prior reproducing test is a §11.4.43/§11.4.123 violation (no falsifiability evidence the fix addresses the root cause vs masks the symptom). Where the recorded issue cannot be reproduced, the deep-research-before-declaring-untestable path (§11.4.123 / §11.4.150) is mandatory before any “cannot reproduce” classification.
3. **A fix / improvement for every confirmed issue.** Every confirmed Crashlytics issue (severity per the §11.4.47 classification table) MUST receive a real fix or improvement landed against its proven root cause (§11.4.9 / §11.4.43 / §11.4.108 four-layer verification on a clean artifact). “Acknowledged in the console” is not a fix; muting / ignoring a recurring issue without a tracked rationale is a §11.4.6 / §11.4.90 violation.
4. **Validation + verification regression-test coverage per closed issue.** Every Crashlytics issue closed by a fix MUST, in the SAME commit as the fix (§11.4.43 DOCUMENT step), register a permanent §11.4.135 regression guard into the standing suite — a §11.4.115 polarity test whose RED_MODE=1 captures the recorded defect on the pre-fix artifact and whose RED_MODE=0 is the standing GREEN guard asserting the defect is ABSENT. The guard MUST exercise the same user-reachable code path the Crashlytics stacktrace / trace identifies, with rock-solid captured evidence per §11.4.5 / §11.4.69 / §11.4.107 / §11.4.123. A Crashlytics issue marked resolved WITHOUT a falsifiable regression test is FORBIDDEN — it is the canonical recurrence vector (the issue silently returns months later, invisible to the green suite), the §11.4.138 operator-escape class applied to field telemetry. Each closure carries the §11.4.34 demotion-evidence + a closure-log entry recording the Crashlytics issue id / console URL, the root-cause analysis, the fix commit, and the validation + verification test paths.

5. Regular cadence — checked repeatedly, never once. The monitoring + debug + fix + cover cycle MUST fire on a regular cadence, reusing the §11.4.47 five-trigger set (pre-build / pre-flash / pre-distribute / pre-tag blocking; daily + post-deployment-burn-in non-blocking) so newly-recorded issues are caught while they are cheap. A one-time sweep that is never repeated is a §11.4.152 violation; field telemetry accrues continuously and so must its processing.

6. No false results, no bluff. Every verdict in the cycle is captured evidence per §11.4 / §11.4.1 / §11.4.6 — a “no new Crashlytics issues” claim requires the enumerated monitored-surfaces + queried-window as proof (§11.4.118 absence-of-evidence is not evidence-of-absence); a “fixed” claim requires the RED→GREEN flip on the artifact (§11.4.115 / §11.4.130); a regression guard whose paired §1.1 mutation does not make it FAIL is itself a bluff gate.

Honest boundary (§11.4.6): processing every recorded issue reduces the known-field-defect surface and breaks the silent-recurrence vector — it does NOT prove zero remaining field defects (a defect no user has hit yet is not yet recorded), and it does NOT replace §11.4.108 runtime-signature verification or §11.4.40 full-suite retest. It is the field-telemetry-driven layer that complements them; an issue is “closed” only when its regression guard is GREEN on a clean artifact, never when the console mark is flipped.

Classification: universal (§11.4.17) — a platform-neutral field-telemetry-processing discipline reusable by ANY project that wires Firebase Crashlytics (or, by the same shape, any equivalent crash/ANR/perf/non-fatal field-reporting backend); the consuming project supplies its concrete Crashlytics project handle, console-access credential (§11.4.10, never logged), severity table, and per-issue closure-log path per §11.4.35. The reference project-level instantiation is the Lava client’s §6.O (Crashlytics-Resolved Issue Coverage Mandate — per-issue validation test + Challenge test + .lava-ci-evidence/crashlytics-resolved/<date>-<slug>.md closure log) + §6.AC (Comprehensive Non-Fatal Telemetry Mandate — every catch / fallback / recovered-error path records a non-fatal with triage context), which together are the project-specific embodiment of this universal clause. Composes §11.4 / §11.4.1 / §11.4.5 / §11.4.6 / §11.4.9 / §11.4.34 / §11.4.40 / §11.4.43 / §11.4.47 / §11.4.69 / §11.4.90 / §11.4.102 / §11.4.107 / §11.4.108 / §11.4.115 / §11.4.118 / §11.4.123 / §11.4.130 / §11.4.135 / §11.4.138 / §11.4.150 / §1.1. Propagation gate CM-COVENANT-114-152-PROPAGATION (literal 11.4.152 across the consumer fleet) + recommended gate CM-CRASHLYTICS-ISSUE-FULLY-COVERED (every closed Crashlytics issue carries a systematic-debug root-cause record + a registered §11.4.135 regression guard + a closure-log entry citing the console issue id/URL + the validation/verification test paths; a console-resolved issue with no falsifiable regression guard FAILs) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; mark a Crashlytics issue resolved with no registered regression guard / no closure log → CM-CRASHLYTICS-ISSUE-FULLY-COVERED FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md §11.4.152. Non-compliance is a release blocker. No escape hatch — no --skip-crashlytics-monitoring, --monitor-crashes-only, --skip-non-fatals, --resolve-without-regression-test, --console-mark-is-fixed, --monitor-once, --mute-without-rationale flag.

§11.4.153 — Comprehensive per-feature Status + Status_Summary document set with mandatory video-recording confirmation (User mandate, 2026-06-15)

Forensic anchor — verbatim user mandate (2026-06-15): “We MUST CREATE proper Status and Status_Summary documents under docs/features containing and covering all system components, all client apps, and all features we have and we have ported from all cli_agents! ... in-depth with no a single feature left ... table ... always up to date and in sync ... status of its usability implementation tests coverage and the confirmation of practical use confirmed on video recording with no false results or bluff of any kind! ... tightly connected using docs_chain Submodule ... Existence of such Status (with Status_Summary) document with all exported files (PDF, DOCX, HTML) is mandatory for every single project.”

Every project MUST maintain, under docs/features/, a comprehensive **feature Status document set** (Status.md + its §11.4.56 Status_Summary.md companion) that enumerates EVERY system component, EVERY client application/binary/surface (TUI / CLI / Web / desktop / mobile / API / gRPC / library / submodule / infrastructure), and EVERY feature — including features ported in from any incorporated CLI-agent / submodule catalogue (§11.4.74) — with NO single feature left out. The mandate (ALL hold):

1. **Total, categorized feature coverage** — a per-feature table covering the FULL feature surface, organized Component → Category, reconciled against the actual codebase (a feature present in code but absent from the table, or a table row with no code, is a §11.4.153 violation; “looks complete” without a code-reconciliation pass is a §11.4.6/§11.4.118 bluff).
2. **Per-feature fields** — each feature row MUST carry: Component, Feature, Category, Implementation status, Wiring status (genuinely reachable by an end user, not merely compiled — §11.4.108), Real-use status, Tests-coverage status (the four-layer §11.4.4(b) coverage state), Validation status (closed vocabulary PASS / FAIL / SKIP / PENDING_FORENSICS / OPERATOR-BLOCKED per §11.4.45), and **Video-recording confirmation** (a path to the captured real-use video proving the feature works for the end user, or an honest gap marker).
3. **Mandatory per-feature real-use video confirmation** — each user-visible feature’s “confirmed working” claim MUST be backed by a recorded real-use video showing a genuine end-user scenario: real prompts → real LLM/service responses producing real results, NEVER a frozen/stale frame (§11.4.107 liveness), NEVER a faked, mocked, demo-loop, or bluff response, NEVER an LLM error/garbage response passed off as success. Videos are captured evidence per §11.4.2/§11.4.5 and stored at the project-declared recording path (§11.4.35); a “confirmed” row with no real video, or a video whose content is a bluff, is a §11.4 PASS-bluff at the feature-status layer. Where autonomous video capture is genuinely infeasible the row is an honest §11.4.3/§11.4.52 SKIP-with-reason + tracked migration item — NEVER a faked confirmation.
4. **Mandatory video-analysis remediation loop** — every recorded video MUST be analysed (its presented data recognized + checked); any defect the video surfaces MUST trigger the §11.4.102 systematic-debugging → fix → §11.4.146 retest → re-record loop, iterating to a clean GO per §11.4.134, before the feature is marked confirmed. A video that exposes a broken feature is a §11.4.4 test-interrupt, not a confirmation.

5. **Always-in-sync, mechanically enforced** — the feature Status set is a §11.4.45-class roster/corpus-backed Status doc: kept in sync via the §11.4.106 docs_chain engine + a §11.4.86 drift-proof fingerprint (sha256 of the sorted feature-key roster AND the sorted video-artefact roster, NOT mtime) so that adding/removing/renaming a feature OR a video artefact mechanically re-syncs the docs + exports out-of-the-box; stale = violation.
6. **Four-format export** — the feature Status set MUST export to HTML + PDF + DOCX (this doc class ADDS DOCX to the §11.4.65 universal HTML+PDF export set; the other doc classes are unchanged), all kept in sync per §11.4.60.
7. **Follows the Status-doc family** — composes §11.4.44 (revision header), §11.4.45 (Status maintenance + captured-evidence table + operator-blocked-at-top), §11.4.56 (Status_Summary two-audience parity), §11.4.57 (README doc-link row), §11.4.59/§11.4.60 (always-sync composite).

Honest boundary (§11.4.6): the feature Status set guarantees a complete, video-confirmed, always-synced ledger of what works for the end user — it does NOT itself prove correctness of any single feature (that rests on the feature's own §11.4.69/§11.4.107/§11.4.123 captured evidence the row cites), and it does NOT replace §11.4.40 full-suite retest. Classification: universal (§11.4.17) — the consuming project supplies its concrete docs/features/ layout, feature roster source, recording path, and DOCX exporter per §11.4.35. Composes §11.4.2 / §11.4.5 / §11.4.44 / §11.4.45 / §11.4.52 / §11.4.56 / §11.4.57 / §11.4.59 / §11.4.60 / §11.4.65 / §11.4.86 / §11.4.102 / §11.4.106 / §11.4.107 / §11.4.108 / §11.4.118 / §11.4.123 / §11.4.134 / §11.4.146 / §1.1. Propagation gate CM-COVENANT-114-153-PROPAGATION (literal 11.4.153) + recommended gates CM-FEATURE-STATUS-COMPLETE (every code-present feature has a Status row + every row maps to code) + CM-FEATURE-STATUS-VIDEO-CONFIRMED (every user-visible confirmed row cites a real captured video or an honest SKIP) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; mark a feature confirmed with no real video / drop a code-present feature from the table → the feature-status gates FAIL; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md §11.4.153. Non-compliance is a release blocker. No escape hatch — no --skip-feature-status, --feature-without-video, --frozen-video-OK, --bluff-response-video-OK, --skip-video-analysis, --feature-ledger-incomplete-OK, --no-docx-export, --allow-feature-ledger-drift flag.

§12. Host-session safety — directly OR indirectly signing the user out is FORBIDDEN

Every script, test, helper, and AI agent governed by this Constitution MUST respect host-session safety. Non-compliance is a release blocker.

§12.1 Forbidden operations — directly OR indirectly

1. **Suspending the host:** systemctl suspend, pm-suspend, loginctl suspend, Dbus org.freedesktop.login1.Suspend, GNOME idle-suspend, lid-close handler.

2. **Hibernating / hybrid-sleeping:** any Hibernate / HybridSleep / SuspendThenHibernate method.
3. **Logging out the user:** `loginctl terminate-session`, `kill -u <user>`, `systemctl --user --kill`, anything that signals `user@<uid>.service`.
4. **Unbounded-memory operations** inside `user@<uid>.service` cgroup. Any single command expected to exceed ~4 GiB RSS MUST be wrapped in a bounded execution scope (e.g. `bounded_run` from a project helper library that wraps `systemd-run --user --scope -p MemoryMax=...`).
5. **Programmatic rkill toggles, lid-switch handlers, or power-button handlers** — these cascade into idle-actions.
6. **Disabling systemd-logind, GDM, or session managers** “to make things faster” — even temporary stops leave the system unable to recover.

§12.2 Required safeguards

Every script in this project that performs heavy work (build, transcription, model inference, large compression, multi-GB git operations) MUST:

1. Source the project’s host-safety helper library at the top.
2. Call its pre-flight check and **abort if it fails**.
3. Wrap any subprocess expected to exceed ~4 GiB RSS in a bounded execution scope so the kernel OOM killer is contained to that scope and cannot escalate to `user.slice`.
4. Cap parallelism (`-j`) to fit available RAM (estimate per-job-peak-RSS and divide into the budget).

§12.3 Container hygiene

Containers (Docker / Podman) the project owns or relies on MUST:

1. Declare an explicit memory limit (`mem_limit` / `--memory` / `MemoryMax`).
2. Set `OOMPolicy=stop` in their `systemd` unit to avoid retry loops.
3. Use exponential-backoff restart policies, never immediate retry.
4. Be clean-slate destroyed and rebuilt after any host crash or session loss so stale lock files don’t keep producing failures.

§12.6 Memory-Budget Ceiling — 60% MAXIMUM

Forensic anchor — direct user mandate:

“First make sure that whatever we do through our procedures related to this project MUST NOT use more than 60% of total system memory! All processes MUST be able to function normally!”

The mandate. Project procedures MUST NOT use more than **60% of total system RAM**. The remaining 40% is reserved for the operator’s other workloads so the host can keep serving them while project work proceeds.

The protections:

1. `HOST_SAFETY_MAX_MEM_PCT` defaults to 60.
2. `HOST_SAFETY_BUDGET_GB` is computed at source-time from `MemTotal × MAX_PCT/100`, in GiB.

3. Bounded execution scopes clamp MemoryMax down to the budget if the caller asks for more.
4. The build script's parallelism is computed as $\min(\text{nproc}, \text{floor}(\text{budget_gb} / \text{per_job_peak_rss_gb}))$.
5. Heavy work MUST be wrapped in a bounded execution scope so the kernel OOM-kills only the scope — `user@<uid>.service` stays alive.

No escape hatch. §12.6 has NO operator-facing override flag. The cap exists for the operator's own protection; bypassing it is the bluff the §11.4 covenant specifically prohibits.

§12.10 Continuation document — sacred invariant

Forensic anchor — direct user mandate:

“during any work we perform, during Phases implementation, debugging and fixing, during ANY effort we have the Continuation document MUST BE maintained and it MUST NOT BE out of sync with current work we are doing! If for any reason we stop our work, we MUST BE able to continue any time, with current work, exactly where we have left off and from any CLI agent or any LLM model we chose! Nothing can be broken or faulty in maintained Continuation document!”

The mandate. A single, canonical, machine-readable handoff document — `docs/CONTINUATION.md` — must always reflect the live state of the project. Any agent (human, Claude Code, Cursor, Aider, Codex, Gemini CLI, any future LLM) must be able to resume work **exactly where the previous session left off** by reading this single file.

Mandatory protections:

1. **`docs/CONTINUATION.md` MUST exist** at the project root. Its absence is a release blocker.
2. **Every non-trivial state change** MUST update this document in the same commit as the work itself.
3. **Top-of-file timestamp** is updated on every edit. Stale timestamps trigger gate failure.
4. **Section §3 “Active work”** lists every IN PROGRESS / BLOCKED item with enough detail that any agent can resume without conversation context — concrete commands, file paths, monitor IDs.
5. **Section §0 “How to use this document”** contains the verbatim resumption prompt — a single block any operator can paste into any CLI agent.
6. **Document MUST be self-contained.** No hyperlinks to ephemeral external systems as the only source of truth.

No escape hatch. §12.10 has NO operator-facing override flag for the existence requirement. The discipline exists for the operator's own protection.

Appendix A — Mutation testing — academic and industrial foundations

The anti-bluff / paired-mutation policy in §1.1 is rooted in the **mutation testing** literature. References that the consuming project SHOULD cite when explaining its meta-test harness:

- Jia, Y. & Harman, M. (2011). *An Analysis and Survey of the Development of Mutation Testing*. IEEE Transactions on Software Engineering, 37(5).
- DeMillo, R.A., Lipton, R.J., Sayward, F.G. (1978). *Hints on Test Data Selection: Help for the Practicing Programmer*. IEEE Computer.
- Open-source mutation testers worth studying:
 - **PIT** (Java) — <https://pitest.org/>
 - **Stryker** (JS / C# / Scala) — <https://stryker-mutator.io/>
 - **Cosmic Ray** (Python) — <https://github.com/sixty-north/cosmic-ray>
 - **mutmut** (Python) — <https://mutmut.readthedocs.io/>
 - **mull** (LLVM-IR) — <https://mull.readthedocs.io/>

The §1.1 paired-mutation policy is a **lightweight** variant of mutation testing applied at the gate-assertion layer rather than at the production-code layer. It does not need a mutation framework; it needs ONE mutation per gate that proves the gate catches the break. This is computationally cheap ($O(\text{gates})$, not $O(\text{production code lines})$) and produces a strong “bluff-immunity” guarantee.

Appendix B — Recursive inheritance & path-independence

Projects that include this constitution submodule MUST tolerate **arbitrary submodule depth**. Between the main project root and any consuming submodule there may be N intermediate levels. To locate this constitution submodule from any depth, every project SHOULD provide a helper that walks up parents until it finds constitution/Constitution.md (the file you are reading right now) OR follows `git rev-parse --show-superproject-working-tree` recursively. The constitution submodule ships `find_constitution.sh` for exactly this purpose; nested submodules can source it without knowing their own depth.

Appendix C — Multi-upstream push

This constitution submodule is hosted on multiple Git providers, and consuming projects often are too. The submodule ships an `install_upstreams.sh` helper (or invokes the `install_upstreams` system command from the parent toolkit) that reads `Upstreams/*.sh` declarations and configures all remotes locally.

Every commit to the constitution submodule MUST be pushed to ALL configured upstreams. Use a shell helper that iterates `git remote` and pushes one by one, OR configure origin with multiple push URLs (`git remote set-url --add --push origin <url>` per remote).
