

APPENDIX

Appendix: Manifests, Schemas, Templates & Full Source Code

Revision: 1 Last modified: 2026-07-24T00:00:00Z

Table of Contents

- [A. Extension Manifest](#)
 - [B. Nano-Task Schema](#)
 - [C. Workable Items Hierarchy Example](#)
 - [D. MCP Server Configuration](#)
 - [E. Helix LLM Gateway](#) — `main.go`
 - [F. SuperBridge MCP Server](#) — `src/index.ts`
 - [G. Implementation Gate](#) — `implementation-gate.js`
 - [H. Task Granulator](#) — `granulator.ts`
 - [I. DAG Engine](#) — `dag.go`
 - [J. Constitution-Aware Commands](#) — [Bash Scripts](#)
 - [K. Pack Distribution Script](#) — `pack.sh`
 - [L. Environment Template](#) — `.env.template`
 - [M. Container Compose Files](#)
 - [N. Test Artifact Template](#)
 - [O. CI Configuration](#) — Disabled per §11.4.156
-

A. Extension Manifest

File: `extension.yml`

extension.yml — Speckit Superpowers Bridge Extension Manifest
§11.4.31: helix-deps.yaml declares own-org dependencies

name: speckit-superpowers-bridge

version: "1.0.0"

description: >

Full Speckit → Superpowers bridge with nano-task decomposition,
TDD-enforced implementation gates, DAG-based dependency resolution,
and Helix LLM-integrated autonomous execution.

author: HelixDevelopment

license: Apache-2.0

constitution:

submodule: constitution

min_revision: 227

required_anchors:

- "11.4.43" # TDD-fix-discipline
- "11.4.115" # RED-baseline + polarity switch
- "11.4.142" # Universal code-review
- "11.4.224" # TDD for ALL work
- "11.4.146" # Reproduce-first-then-extend
- "11.4.102" # Systematic-debugging auto-activation
- "11.4.108" # Four-layer verification
- "11.4.69" # Captured-evidence taxonomy
- "11.4.161" # Rootless containers
- "11.4.156" # CI disabled

components:

- name: superbridge-mcp

type: mcp-server

language: typescript

entry: src/index.ts

runtime: node >=20

description: MCP server exposing Superpowers skills via Model Context Protocol

- name: implementation-gate

type: tool

language: typescript

entry: src/implementation-gate.ts

runtime: node >=20

description: Mechanical TDD enforcement gate (RED → GREEN → REFACTOR)

- name: task-granulator

type: tool

language: typescript
entry: src/granulator.ts
runtime: node >=20
description: Nano-task decomposition engine

- name: helix-llm-gateway
type: service
language: go
entry: cmd/gateway/main.go
runtime: go >=1.22
description: Helix LLM Gateway — routing, auth, context budgeting

- name: dag-engine
type: library
language: go
entry: pkg/dag/dag.go
runtime: go >=1.22
description: DAG-based nano-task dependency resolution + execution ordering

- name: constitution-commands
type: scripts
language: bash
entry: scripts/
description: Constitution-aware CLI helpers

skills:

- speckit-superpowers-bridge-handoff
- speckit-superpowers-bridge-execute
- speckit-superpowers-bridge-guard

mcp_tools:

- name: tdd_enforce
description: Enforce TDD (RED → GREEN → REFACTOR) for a nano-task
server: superbridge-mcp
- name: verify_anti_bluff
description: Verify a nano-task's TDD verdict is anti-bluff compliant
server: superbridge-mcp
- name: granulate_task
description: Decompose a feature task into nano-tasks
server: superbridge-mcp
- name: resolve_dag
description: Resolve nano-task dependency DAG and return execution order
server: superbridge-mcp
- name: submit_inference
description: Submit a nano-task for LLM inference via Helix Gateway
server: superbridge-mcp

directories:

verdicts: qa-results/tdd_verdicts

evidence: qa-results/tdd_evidence

sandboxes: /tmp/tdd_sandboxes

containers: containers/

dependencies:

- name: vasic-digital/containers

type: submodule

path: submodules/containers

reason: Container orchestration for nano-task sandboxes

layout: flat

- name: vasic-digital/Challenges

type: submodule

path: submodules/challenges

reason: Anti-bluff Challenge banks for TDD enforcement

layout: flat

- name: HelixDevelopment/HelixQA

type: submodule

path: submodules/helix_qa

reason: Autonomous QA sessions for bridge validation

layout: flat

B. Nano-Task Schema

File: schemas/nano-task.schema.json

```

{
  "$schema": "https://json-schema.org/draft/2020-12/schema",
  "$id": "https://speckit-bridge.helix/schemas/nano-task.schema.json",
  "title": "NanoTask",
  "description": "Schema for a single decoupled nano-task in the SpecKit-Superpowers bridge.",
  "type": "object",
  "required": [
    "id",
    "parent_feature_id",
    "title",
    "description",
    "language",
    "test_file",
    "impl_file",
    "layer",
    "dependencies",
    "acceptance_criteria"
  ],
  "properties": {
    "id": {
      "type": "string",
      "pattern": "^(ATM|SPK)-\\d{3,}-NT-\\d{3,}$",
      "description": "Stable nano-task identifier: <parent>-NT-<sequence>"
    },
    "parent_feature_id": {
      "type": "string",
      "pattern": "^(ATM|SPK)-\\d{3,}$",
      "description": "Parent workable item identifier"
    },
    "title": {
      "type": "string",
      "minLength": 10,
      "maxLength": 200,
      "description": "Human-readable nano-task title"
    },
    "description": {
      "type": "string",
      "minLength": 50,
      "description": "Comprehensive description per §11.4.171: what, why, how, who, outcome"
    },
    "language": {
      "type": "string",
      "enum": ["typescript", "javascript", "go", "bash", "yaml", "json", "markdown"],
      "description": "Implementation language"
    }
  }
}

```

```

"test_file": {
  "type": "string",
  "description": "Relative path to the test file (must be authored FIRST per §11.4.224)"
},
"impl_file": {
  "type": "string",
  "description": "Relative path to the implementation file (created AFTER RED phase)"
},
"layer": {
  "type": "integer",
  "minimum": 0,
  "description": "Decomposition layer depth (0 = leaf, N = N layers of sub-tasks)"
},
"dependencies": {
  "type": "array",
  "items": {
    "type": "string",
    "pattern": "^(ATM|SPK)-\\d{3,}-NT-\\d{3,}$"
  },
  "description": "Nano-task IDs this task depends on (must be complete before this can start)"
},
"acceptance_criteria": {
  "type": "array",
  "items": {
    "type": "object",
    "required": ["description", "evidence_class"],
    "properties": {
      "description": {
        "type": "string",
        "minLength": 10
      },
      "evidence_class": {
        "type": "string",
        "enum": ["source", "artifact", "runtime", "user_visible"],
        "description": "Evidence class required for closure per §11.4.226"
      }
    }
  },
  "minItems": 1
},
"estimated_tokens": {
  "type": "integer",
  "minimum": 1,
  "maximum": 4096,
  "description": "Estimated prompt token count for LLM inference"
},

```

```

"context_budget": {
  "type": "integer",
  "maximum": 4096,
  "default": 2048,
  "description": "Maximum context tokens for this nano-task's LLM inference"
},
"timeout_seconds": {
  "type": "integer",
  "default": 300,
  "description": "Hard timeout for RED/GREEN phase execution"
},
"max_cycles": {
  "type": "integer",
  "default": 5,
  "description": "Maximum RED → GREEN cycles before loop detection fires"
},
"sandbox_profile": {
  "type": "string",
  "enum": ["default", "networked", "gpu", "privileged"],
  "default": "default",
  "description": "Container sandbox security profile"
},
"red_mode_env": {
  "type": "object",
  "additionalProperties": { "type": "string" },
  "description": "Environment variables to set for RED phase execution"
},
"green_mode_env": {
  "type": "object",
  "additionalProperties": { "type": "string" },
  "description": "Environment variables to set for GREEN phase execution"
},
"tags": {
  "type": "array",
  "items": { "type": "string" },
  "description": "Categorization tags"
},
"created_at": {
  "type": "string",
  "format": "date-time"
},
"created_by": {
  "type": "string",
  "description": "Canonical handle of creator per §11.4.104"
},
"assigned_to": {

```

```
"type": "string",  
"description": "Canonical handle of assignee per §11.4.104"  
}  
}  
}
```

C. Workable Items Hierarchy Example

File: docs/design/hierarchy_example.yml

Workable Items Hierarchy — N-Layer Decomposition Example
Demonstrates the SpecKit → Superpowers → Nano-Task decomposition chain.

feature:

id: "ATM-500"
title: "User authentication with OAuth 2.0"
type: Feature
status: In progress
spec_kit_phase: implement

Layer 1: SpecKit tasks

tasks:

- id: "ATM-500-T1"
title: "Implement OAuth token exchange endpoint"
type: Task
language: go
spec: "POST /auth/token accepting authorization_code, returning access_token + refresh_token"
dependencies: []
- id: "ATM-500-T2"
title: "Implement token validation middleware"
type: Task
language: go
spec: "Middleware that validates Bearer token, injects user context"
dependencies: ["ATM-500-T1"]
- id: "ATM-500-T3"
title: "Implement session persistence"
type: Task
language: go
spec: "Store refresh tokens in Redis with TTL"
dependencies: ["ATM-500-T1"]

Layer 2: Superpowers-decomposed nano-tasks for T1

nano_tasks:

- id: "ATM-500-T1-NT-001"
parent: "ATM-500-T1"
title: "Define OAuth token exchange request/response types"
layer: 2
language: go
test_file: "pkg/auth/token_types_test.go"
impl_file: "pkg/auth/token_types.go"
dependencies: []
acceptance_criteria:
 - description: "TokenRequest struct parses valid JSON"

```
evidence_class: source
- description: "TokenResponse struct serializes to valid JSON"
  evidence_class: source

- id: "ATM-500-T1-NT-002"
  parent: "ATM-500-T1"
  title: "Implement authorization_code → token exchange logic"
  layer: 2
  language: go
  test_file: "pkg/auth/token_exchange_test.go"
  impl_file: "pkg/auth/token_exchange.go"
  dependencies: ["ATM-500-T1-NT-001"]
  acceptance_criteria:
    - description: "Valid authorization_code returns access_token"
      evidence_class: runtime
    - description: "Expired authorization_code returns 401"
      evidence_class: runtime
    - description: "Invalid grant_type returns 400"
      evidence_class: runtime


- id: "ATM-500-T1-NT-003"
  parent: "ATM-500-T1"
  title: "Implement JWT access token signing"
  layer: 2
  language: go
  test_file: "pkg/auth/jwt_signing_test.go"
  impl_file: "pkg/auth/jwt_signing.go"
  dependencies: ["ATM-500-T1-NT-001"]
  acceptance_criteria:
    - description: "Generated JWT contains sub, exp, iat claims"
      evidence_class: source
    - description: "JWT signature verifies with known public key"
      evidence_class: runtime

- id: "ATM-500-T1-NT-004"
  parent: "ATM-500-T1"
  title: "Wire token exchange endpoint into HTTP router"
  layer: 2
  language: go
  test_file: "pkg/auth/router_wire_test.go"
  impl_file: "pkg/auth/router_wire.go"
  dependencies: ["ATM-500-T1-NT-002", "ATM-500-T1-NT-003"]
  acceptance_criteria:
    - description: "POST /auth/token routes to exchange handler"
      evidence_class: runtime
    - description: "Unknown routes return 404"
```

evidence_class: runtime

Visualization of the dependency DAG for T1:

dag_ascii: |



D. MCP Server Configuration

File: mcp.json

```

{
  "mcpServers": {
    "superbridge": {
      "command": "node",
      "args": [
        "dist/index.js"
      ],
      "cwd": "${PROJECT_ROOT}/superbridge-mcp",
      "env": {
        "HELIX_GATEWAY_URL": "https://localhost:3101",
        "TDD_VERDICT_DIR": "qa-results/tdd_verdicts",
        "TDD_EVIDENCE_DIR": "qa-results/tdd_evidence",
        "TDD_ITERATIONS": "3",
        "TDD_COVERAGE_THRESHOLD": "0.85",
        "TDD_MAX_CYCLES": "5",
        "TDD_SANDBOX_ENABLED": "true",
        "LLAMA_RPC_HOST": "localhost",
        "LLAMA_RPC_PORT": "8080",
        "LLAMA_RPC_TLS_ENABLED": "true",
        "GATEWAY_CLIENT_CERT": "${PROJECT_ROOT}/certs/gateway-client.crt",
        "GATEWAY_CLIENT_KEY": "${PROJECT_ROOT}/certs/gateway-client.key",
        "GATEWAY_CA_CERT": "${PROJECT_ROOT}/certs/ca.crt",
        "RED_MODE": "1"
      },
      "disabled": false
    },
    "codegraph": {
      "command": "npx",
      "args": [
        "@colbymchenry/codegraph",
        "serve",
        "--mcp"
      ],
      "disabled": false
    }
  }
}

```

E. Helix LLM Gateway – `main.go`

File: `cmd/gateway/main.go`

```
package main
```

```
import (  
    "context"  
    "crypto/sha256"  
    "crypto/tls"  
    "crypto/x509"  
    "encoding/hex"  
    "encoding/json"  
    "errors"  
    "flag"  
    "fmt"  
    "io"  
    "log"  
    "net/http"  
    "os"  
    "os/signal"  
    "sync"  
    "syscall"  
    "time"  
)
```

```
// — Configuration
```

```
type Config struct {  
    ListenAddr    string    `json:"listen_addr"`  
    MaxTokens     int       `json:"max_tokens"`  
    RequestTimeout time.Duration `json:"request_timeout"`  
    Backends      []Backend `json:"backends"`  
    TLSCertFile   string    `json:"tls_cert_file"`  
    TLSKeyFile    string    `json:"tls_key_file"`  
    TLSClientCAFile string    `json:"tls_client_ca_file"`  
    ModelManifest string    `json:"model_manifest"`  
}
```

```
type Backend struct {  
    Name    string `json:"name"`  
    URL     string `json:"url"`  
    Priority int    `json:"priority"` // 1 = native (highest), 2+ = provider  
    MaxQueue int    `json:"max_queue"`  
}
```

```
type ModelManifest struct {  
    Models []ModelEntry `json:"models"`  
}
```

```

}

type ModelEntry struct {
    Name    string `json:"name"`
    Path    string `json:"path"`
    SHA256  string `json:"sha256"`
    Source  string `json:"source"`
    Default bool   `json:"default"`
}

```

// — Gateway

```

type Gateway struct {
    config    Config
    manifest  ModelManifest
    server    *http.Server
    backendsMu sync.RWMutex
    rateLimits map[string]*RateLimiter
}

```

```

type RateLimiter struct {
    lastRequest time.Time
    mu          sync.Mutex
}

```

// — Inference Request/Response

```

type InferenceRequest struct {
    TaskID    string `json:"task_id"`
    Prompt    string `json:"prompt"`
    MaxTokens int    `json:"max_tokens"`
    Temperature float64 `json:"temperature"`
    Model     string `json:"model,omitempty"`
}

```

```

type InferenceResponse struct {
    TaskID    string `json:"task_id"`
    Text      string `json:"text"`
    TokensUsed int    `json:"tokens_used"`
    Model     string `json:"model"`
    Backend   string `json:"backend"`
    FinishReason string `json:"finish_reason"`
}

```

// — Constructor

```
func NewGateway(cfg Config) (*Gateway, error) {  
    gw := &Gateway{  
        config:    cfg,  
        rateLimits: make(map[string]*RateLimiter),  
    }  
  
    if cfg.ModelManifest != "" {  
        data, err := os.ReadFile(cfg.ModelManifest)  
        if err != nil {  
            return nil, fmt.Errorf("read model manifest: %w", err)  
        }  
        if err := json.Unmarshal(data, &gw.manifest); err != nil {  
            return nil, fmt.Errorf("parse model manifest: %w", err)  
        }  
    }  
  
    return gw, nil  
}
```

// — Model Integrity Verification

```
func (g *Gateway) verifyModelIntegrity(modelPath, expectedHash string) error {  
    f, err := os.Open(modelPath)  
    if err != nil {  
        return fmt.Errorf("open model: %w", err)  
    }  
    defer f.Close()  
  
    h := sha256.New()  
    if _, err := io.Copy(h, f); err != nil {  
        return fmt.Errorf("hash model: %w", err)  
    }  
    actual := hex.EncodeToString(h.Sum(nil))  
  
    if actual != expectedHash {  
        return fmt.Errorf("MODEL_INTEGRITY_FAIL: %s hash=%s expected=%s",  
            modelPath, actual, expectedHash)  
    }  
    return nil  
}
```

// — Context Budget Enforcement

```
func (g *Gateway) enforceBudget(req InferenceRequest) error {  
    if req.MaxTokens <= 0 {  
        req.MaxTokens = g.config.MaxTokens  
    }  
    if req.MaxTokens > g.config.MaxTokens {  
        return fmt.Errorf("REJECTED: requested %d tokens exceeds budget %d",  
            req.MaxTokens, g.config.MaxTokens)  
    }  
    return nil  
}
```

// — Prompt Sanitization

```
var credentialPatterns = []string{  
    `api[_]?key`,  
    `apikey`,  
    `access[_]?token`,  
    `secret`,  
    `password`,  
    `passwd`,  
}  
  
func sanitizePrompt(prompt string) (string, bool) {  
    sanitized := prompt  
    redacted := false  
    for _, pattern := range credentialPatterns {  
        // In production, use a proper regexp-based scanner  
        if containsCredentialPattern(sanitized, pattern) {  
            sanitized = redactPattern(sanitized, pattern)  
            redacted = true  
        }  
    }  
    return sanitized, redacted  
}  
  
func containsCredentialPattern(s, pattern string) bool {  
    return len(s) > 0 && len(pattern) > 0 // Stub — real impl uses regexp  
}  
  
func redactPattern(s, pattern string) string {  
    return s // Stub — real impl redacts matched text  
}
```

// — Backend Selection (§11.4.196 native-first) —

```
func (g *Gateway) selectBackend() (*Backend, error) {
    g.backendsMu.RLock()
    defer g.backendsMu.RUnlock()

    // Native-first: scan all native backends FIRST (§11.4.196(A))
    var bestBackend *Backend
    for i := range g.config.Backends {
        b := &g.config.Backends[i]
        if b.Priority == 1 { // Native
            if g.isOperational(b) {
                bestBackend = b
                break
            }
        }
    }

    // Fallback to providers if no native is operational
    if bestBackend == nil {
        for i := range g.config.Backends {
            b := &g.config.Backends[i]
            if b.Priority > 1 && g.isOperational(b) {
                bestBackend = b
                break
            }
        }
    }

    if bestBackend == nil {
        return nil, errors.New("no operational backend available — all are rate-limited or unreachable")
    }

    return bestBackend, nil
}

func (g *Gateway) isOperational(b *Backend) bool {
    rl, ok := g.rateLimits[b.Name]
    if !ok {
        return true
    }
    rl.mu.Lock()
    defer rl.mu.Unlock()

    // Rate-limited if last request was within cooldown window
    cooldown := 5 * time.Second
```

```

    return time.Since(rl.lastRequest) > cooldown
}

// — Inference Handler

func (g *Gateway) handleInference(w http.ResponseWriter, r *http.Request) {
    if r.Method != http.MethodPost {
        http.Error(w, "Method not allowed", http.StatusMethodNotAllowed)
        return
    }

    body, err := io.ReadAll(io.LimitReader(r.Body, 1<<20)) // 1MB limit
    if err != nil {
        http.Error(w, "Failed to read request body", http.StatusBadRequest)
        return
    }
    defer r.Body.Close()

    var req InferenceRequest
    if err := json.Unmarshal(body, &req); err != nil {
        writeJSON(w, http.StatusBadRequest, map[string]string{
            "error": fmt.Sprintf("Invalid JSON: %v", err),
        })
        return
    }

    // Enforce context budget
    if err := g.enforceBudget(req); err != nil {
        writeJSON(w, http.StatusPaymentRequired, map[string]string{"error": err.Error()})
        return
    }

    // Sanitize prompt for credentials
    sanitized, redacted := sanitizePrompt(req.Prompt)
    if redacted {
        log.Printf("WARNING: Credential pattern redacted in task %s", req.TaskID)
    }
    req.Prompt = sanitized

    // Select backend
    backend, err := g.selectBackend()
    if err != nil {
        writeJSON(w, http.StatusServiceUnavailable, map[string]string{"error": err.Error()})
        return
    }

```

```

// Forward to backend (simplified — real impl calls llama.cpp RPC)
ctx, cancel := context.WithTimeout(r.Context(), g.config.RequestTimeout)
defer cancel()

resp := g.forwardInference(ctx, backend, req)

// Log audit event
g.logAuditEvent(req.TaskID, backend.Name, resp.TokensUsed)

writeJSON(w, http.StatusOK, resp)
}

func (g *Gateway) forwardInference(ctx context.Context, backend *Backend, req InferenceRequest)
    InferenceResponse {
    // In production: actually forward to llama.cpp RPC or cloud provider
    // This is a stub returning a placeholder response
    select {
    case <-ctx.Done():
        return InferenceResponse{
            TaskID:    req.TaskID,
            FinishReason: "timeout",
        }
    default:
        return InferenceResponse{
            TaskID:    req.TaskID,
            Text:      "[stub inference response]",
            TokensUsed: len(req.Prompt) / 4,
            Model:     "local-model",
            Backend:   backend.Name,
            FinishReason: "stop",
        }
    }
}

// — Health Check

func (g *Gateway) handleHealth(w http.ResponseWriter, r *http.Request) {
    backends := make([]map[string]interface{}, 0, len(g.config.Backends))
    for _, b := range g.config.Backends {
        backends = append(backends, map[string]interface{}{
            "name":    b.Name,
            "operational": g.isOperational(&b),
        })
    }
}

```

```

writeJSON(w, http.StatusOK, map[string]interface{}{
    "status": "healthy",
    "backends": backends,
    "budget": map[string]int{
        "max_tokens": g.config.MaxTokens,
    },
})
}

```

// — Audit Logging

```

func (g *Gateway) logAuditEvent(taskID, backend string, tokensUsed int) {
    entry := map[string]interface{}{
        "event":    "inference_complete",
        "task_id":  taskID,
        "backend":  backend,
        "tokens_used": tokensUsed,
        "timestamp": time.Now().UTC().Format(time.RFC3339),
    }
    data, _ := json.Marshal(entry)
    log.Printf("AUDIT: %s", string(data))
}

```

// — Server Startup

```

func (g *Gateway) Start() error {
    mux := http.NewServeMux()
    mux.HandleFunc("/v1/inference", g.handleInference)
    mux.HandleFunc("/health", g.handleHealth)

    g.server = &http.Server{
        Addr:      g.config.ListenAddr,
        Handler:   mux,
        ReadTimeout: 10 * time.Second,
        WriteTimeout: 30 * time.Second,
        IdleTimeout: 120 * time.Second,
    }

    if g.config.TLSCertFile != "" && g.config.TLSKeyFile != "" {
        tlsConfig := &tls.Config{
            MinVersion: tls.VersionTLS13,
        }
    }
}

```

```

    if g.config.TLSClientCAFile != "" {
        caCert, err := os.ReadFile(g.config.TLSClientCAFile)
        if err != nil {
            return fmt.Errorf("read CA cert: %w", err)
        }
        caCertPool := x509.NewCertPool()
        caCertPool.AppendCertsFromPEM(caCert)
        tlsConfig.ClientAuth = tls.RequireAndVerifyClientCert
        tlsConfig.ClientCAs = caCertPool
    }

    g.server.TLSConfig = tlsConfig
    log.Printf("Starting Helix LLM Gateway on %s (mTLS)", g.config.ListenAddr)
    return g.server.ListenAndServeTLS(g.config.TLSCertFile, g.config.TLSKeyFile)
}

log.Printf("Starting Helix LLM Gateway on %s (plain HTTP — development only)",
    g.config.ListenAddr)
return g.server.ListenAndServe()
}

func (g *Gateway) Shutdown(ctx context.Context) error {
    return g.server.Shutdown(ctx)
}

// ——— Helpers

func writeJSON(w http.ResponseWriter, status int, data interface{}) {
    w.Header().Set("Content-Type", "application/json")
    w.WriteHeader(status)
    json.NewEncoder(w).Encode(data)
}

// ——— Main

func main() {
    configPath := flag.String("config", "config/gateway.json", "Path to gateway config")
    flag.Parse()

    data, err := os.ReadFile(*configPath)
    if err != nil {
        log.Fatalf("Failed to read config: %v", err)
    }

```

```

}

var cfg Config
if err := json.Unmarshal(data, &cfg); err != nil {
    log.Fatalf("Failed to parse config: %v", err)
}

gw, err := NewGateway(cfg)
if err != nil {
    log.Fatalf("Failed to create gateway: %v", err)
}

go func() {
    if err := gw.Start(); err != nil && err != http.ErrServerClosed {
        log.Fatalf("Gateway failed: %v", err)
    }
}()

quit := make(chan os.Signal, 1)
signal.Notify(quit, syscall.SIGINT, syscall.SIGTERM)
<-quit

log.Println("Shutting down gateway...")
ctx, cancel := context.WithTimeout(context.Background(), 30*time.Second)
defer cancel()

if err := gw.Shutdown(ctx); err != nil {
    log.Fatalf("Gateway forced shutdown: %v", err)
}
log.Println("Gateway stopped.")
}

```

F. SuperBridge MCP Server – `src/index.ts`

File: `src/index.ts`

```
#!/usr/bin/env node
```

```
/**
```

```
 * SuperBridge MCP Server
```

```
 *
```

```
 * Exposes SpecKit-Superpowers bridge capabilities via the
```

```
 * Model Context Protocol (MCP). Provides TDD enforcement,
```

```
 * anti-bluff verification, task granulation, and DAG resolution.
```

```
 *
```

```
 * Constitutional anchors: §11.4.43, §11.4.115, §11.4.146, §11.4.224
```

```
 */
```

```
import { Server } from '@modelcontextprotocol/sdk/server/index.js';
```

```
import { StdioServerTransport } from '@modelcontextprotocol/sdk/server/stdio.js';
```

```
import {
```

```
  CallToolRequestSchema,
```

```
  ListToolsRequestSchema,
```

```
  Tool,
```

```
} from '@modelcontextprotocol/sdk/types.js';
```

```
import { execSync } from 'child_process';
```

```
import { readFileSync, existsSync, statSync, writeFileSync, mkdirSync } from 'fs';
```

```
import { join, dirname } from 'path';
```

```
import { createHash } from 'crypto';
```

```
// ——— Tool Definitions
```

```
const TOOLS: Tool[] = [
```

```
 {
```

```
   name: 'tdd_enforce',
```

```
   description: 'Enforce TDD (RED → GREEN → REFACTOR) for a nano-task per §11.4.43/  
                §11.4.115',
```

```
   inputSchema: {
```

```
     type: 'object',
```

```
     properties: {
```

```
       command: {
```

```
         type: 'string',
```

```
         enum: ['red', 'green', 'refactor'],
```

```
         description: 'TDD phase to enforce'
```

```
     },
```

```
     item_id: {
```

```
       type: 'string',
```

```
       description: 'Nano-task identifier (ATM-NNN-NT-NNN)'
```

```
     },
```

```
     test_files: {
```

```
       type: 'array',
```

```

    items: { type: 'string' },
    description: 'Paths to test files'
  },
  impl_files: {
    type: 'array',
    items: { type: 'string' },
    description: 'Paths to implementation files'
  }
},
required: ['command', 'item_id', 'test_files']
}
},
{
  name: 'verify_anti_bluff',
  description: 'Verify a nano-task TDD verdict is anti-bluff compliant per §11.4.146(D3)',
  inputSchema: {
    type: 'object',
    properties: {
      item_id: { type: 'string', description: 'Nano-task identifier' },
      red_verdict: { type: 'string', description: 'Path to RED verdict JSON' },
      green_verdict: { type: 'string', description: 'Path to GREEN verdict JSON' }
    },
    required: ['item_id', 'red_verdict', 'green_verdict']
  }
},
{
  name: 'granulate_task',
  description: 'Decompose a feature task into nano-tasks with dependency resolution',
  inputSchema: {
    type: 'object',
    properties: {
      parent_id: { type: 'string', description: 'Parent workable item identifier' },
      spec: { type: 'string', description: 'Feature specification text' },
      language: { type: 'string', description: 'Target implementation language' },
      max_layer_depth: { type: 'integer', default: 3, description: 'Maximum decomposition depth' }
    },
    required: ['parent_id', 'spec', 'language']
  }
},
{
  name: 'resolve_dag',
  description: 'Resolve nano-task dependency DAG and return topological execution order',
  inputSchema: {
    type: 'object',
    properties: {
      parent_id: { type: 'string', description: 'Parent feature identifier' },

```

```

    nano_tasks: {
      type: 'array',
      items: {
        type: 'object',
        properties: {
          id: { type: 'string' },
          dependencies: { type: 'array', items: { type: 'string' } }
        },
        required: ['id']
      }
    },
    required: ['parent_id', 'nano_tasks']
  }
},
{
  name: 'submit_inference',
  description: 'Submit a nano-task for LLM inference via Helix Gateway',
  inputSchema: {
    type: 'object',
    properties: {
      task_id: { type: 'string' },
      prompt: { type: 'string' },
      max_tokens: { type: 'integer', default: 2048 }
    },
    required: ['task_id', 'prompt']
  }
}
];

```

// — Server

```

const server = new Server(
  { name: 'superbridge-mcp', version: '1.0.0' },
  { capabilities: { tools: {} } }
);

```

// — Tool Handlers

```

server.setRequestHandler(ListToolsRequestSchema, async () => ({ tools: TOOLS }));

```

```

server.setRequestHandler(CallToolRequestSchema, async (request) => {
  const { name, arguments: args } = request.params;

```

```

const itemId = (args as any).item_id || (args as any).task_id || "";
const evidenceDir = join('qa-results', 'tdd_evidence', itemId);

try {
  switch (name) {
    case 'tdd_enforce': {
      const { command, test_files, impl_files } = args as any;
      const gateScript = join(process.cwd(), 'node_modules/.bin/implementation-gate');

      // Pass through to implementation-gate.js
      const files = [...(test_files || []), ...(impl_files || [])];
      const cmd = `node ${gateScript} ${command} ${itemId} ${files.join(' ')}`;

      try {
        const output = execSync(cmd, { encoding: 'utf-8', stdio: 'pipe' });
        return {
          content: [{ type: 'text', text: `PASS: ${command.toUpperCase()} phase for ${itemId}\n${output}` }],
        };
      } catch (err: any) {
        return {
          content: [{ type: 'text', text: `REJECTED: ${command.toUpperCase()} phase for ${itemId}\n${err.stdout || ""}\n${err.stderr || ""}` }],
          isError: true
        };
      }
    }

    case 'verify_anti_bluff': {
      const { red_verdict, green_verdict } = args as any;

      if (!existsSync(red_verdict)) {
        return {
          content: [{ type: 'text', text: `FAIL: RED verdict not found at ${red_verdict}` }],
          isError: true
        };
      }

      if (!existsSync(green_verdict)) {
        return {
          content: [{ type: 'text', text: `FAIL: GREEN verdict not found at ${green_verdict}` }],
          isError: true
        };
      }

      const red = JSON.parse(readFileSync(red_verdict, 'utf-8'));
      const green = JSON.parse(readFileSync(green_verdict, 'utf-8'));
    }
  }
}

```

```

const findings: string[] = [];

if (red.exit_code === 0) findings.push('RED_BLUFF: RED passed — tautology');
if (green.exit_code !== 0) findings.push(`GREEN_FAILED: GREEN exited ${green.exit_code}`);
if (red.target_fingerprint === green.target_fingerprint) {
  findings.push('GREEN_BLUFF: Identical fingerprints — fix never deployed');
}
if (red.iterations < 3) findings.push('NONDETERMINISTIC: RED <3 iters');
if (green.iterations < 3) findings.push('NONDETERMINISTIC: GREEN <3 iters');
if (!green.prior_red_verdict) findings.push('MISSING_RED_REF');

for (const f of [...(red.evidence_files || []), ...(green.evidence_files || [])]) {
  if (!existsSync(f)) {
    findings.push(`MISSING_EVIDENCE: ${f}`);
  } else if (statSync(f).size === 0) {
    findings.push(`EMPTY_EVIDENCE: ${f}`);
  }
}

if (findings.length === 0) {
  return {
    content: [{ type: 'text', text: `PASS: ${itemId} — anti-bluff verified (RED → GREEN, distinct
      fingerprints, ≥3 iters, evidence present) [evidence: ${evidenceDir}]` }],
  };
}

return {
  content: [{ type: 'text', text: `FAIL: ${itemId} — ${findings.length} bluff(s):\n${findings.map(f =>
    ` - ${f}`).join('\n')}` }],
  isError: true
};
}

case 'granulate_task': {
  const { parent_id, spec } = args as any;

  // Decompose spec into nano-tasks
  const nanoTasks = decomposeSpec(spec, parent_id, (args as any).max_layer_depth || 3);
  const dag = buildDAG(nanoTasks);
  const order = topologicalSort(dag);

  const result = {
    parent_id,
    nano_task_count: nanoTasks.length,
    max_layer: Math.max(...nanoTasks.map(n => n.layer)),
    execution_order: order,
    nano_tasks: nanoTasks
  };
}

```

```

const resultPath = join(evidenceDir, `GRANULATE_${parent_id}.json`);
mkdirSync(dirname(resultPath), { recursive: true });
writeFileSync(resultPath, JSON.stringify(result, null, 2));

return {
  content: [{ type: 'text', text: `Granulated ${parent_id} into ${nanoTasks.length} nano-tasks
    across ${result.max_layer} layers. Execution order: ${order.join(' → ')} [evidence: $
    {resultPath}]` }],
};
}

case 'resolve_dag': {
  const { nano_tasks } = args as any;
  const dag = buildDAG(nano_tasks);
  const order = topologicalSort(dag);

  return {
    content: [{
      type: 'text',
      text: `DAG resolved for ${args as any}.parent_id: ${nano_tasks.length} nodes, execution
        order: ${order.join(' → ')}
    ]
  ];
}

case 'submit_inference': {
  const { prompt, max_tokens } = args as any;
  const gatewayURL = process.env.HELIX_GATEWAY_URL || 'https://localhost:3101';

  // Forward to Helix LLM Gateway (in production: real HTTP call)
  const response = {
    task_id: itemId,
    text: '[inference result via gateway]',
    tokens_used: Math.floor(prompt.length / 4),
    model: 'local-model',
    backend: 'llama.cpp-native',
    finish_reason: 'stop'
  };

  return {
    content: [{ type: 'text', text: `Inference complete for ${itemId}: ${response.tokens_used} tokens,
      backend=${response.backend}` }],
  ];
}

default:

```

```

    return {
      content: [{ type: 'text', text: `Unknown tool: ${name}` }],
      isError: true
    };
  }
} catch (err: any) {
  return {
    content: [{ type: 'text', text: `Error: ${err.message}` }],
    isError: true
  };
}
});

```

// ——— DAG Implementation

```

interface NanoTaskNode {
  id: string;
  dependencies: string[];
  layer: number;
  title: string;
}

```

```

interface DAG {
  nodes: Map<string, NanoTaskNode>;
  adjacency: Map<string, Set<string>>;
}

```

```

function buildDAG(tasks: NanoTaskNode[]): DAG {
  const dag: DAG = { nodes: new Map(), adjacency: new Map() };

```

```

  for (const task of tasks) {
    dag.nodes.set(task.id, task);
    if (!dag.adjacency.has(task.id)) {
      dag.adjacency.set(task.id, new Set());
    }
    for (const dep of (task.dependencies || [])) {
      if (!dag.adjacency.has(dep)) {
        dag.adjacency.set(dep, new Set());
      }
      dag.adjacency.get(dep)!.add(task.id);
    }
  }
}

```

// Cyclic check

```

const visiting = new Set<string>();

```

```

const visited = new Set<string>();

function dfs(node: string) {
  if (visiting.has(node)) throw new Error(`CYCLE_DETECTED: ${node}`);
  if (visited.has(node)) return;
  visiting.add(node);
  for (const neighbor of dag.adjacency.get(node) || []) {
    dfs(neighbor);
  }
  visiting.delete(node);
  visited.add(node);
}

for (const node of dag.nodes.keys()) {
  if (!visited.has(node)) dfs(node);
}

return dag;
}

function topologicalSort(dag: DAG): string[] {
  const inDegree = new Map<string, number>();
  const queue: string[] = [];
  const result: string[] = [];

  for (const node of dag.nodes.keys()) {
    inDegree.set(node, 0);
  }

  for (const [from, tos] of dag.adjacency) {
    if (!dag.nodes.has(from)) continue;
    for (const to of tos) {
      inDegree.set(to, (inDegree.get(to) || 0) + 1);
    }
  }

  for (const [node, degree] of inDegree) {
    if (dag.nodes.has(node) && degree === 0) {
      queue.push(node);
    }
  }

  while (queue.length > 0) {
    const node = queue.shift()!;
    result.push(node);
    for (const neighbor of dag.adjacency.get(node) || []) {
      const newDegree = (inDegree.get(neighbor) || 1) - 1;

```

```

    inDegree.set(neighbor, newDegree);
    if (newDegree === 0 && dag.nodes.has(neighbor)) {
      queue.push(neighbor);
    }
  }
}

if (result.length !== dag.nodes.size) {
  throw new Error(`DAG_RESOLVE_FAILED: Only ${result.length}/${dag.nodes.size} nodes
    sortable — cycle detected`);
}

return result;
}

function decomposeSpec(spec: string, parentId: string, maxDepth: number): NanoTaskNode[] {
  // In production: invoke LLM (via Helix Gateway) to decompose the spec into nano-tasks.
  // This stub returns a sample decomposition.
  const tasks: NanoTaskNode[] = [
    {
      id: `${parentId}-NT-001`,
      title: 'Define types and interfaces',
      layer: 0,
      dependencies: []
    },
    {
      id: `${parentId}-NT-002`,
      title: 'Implement core logic',
      layer: 0,
      dependencies: [`${parentId}-NT-001`]
    },
    {
      id: `${parentId}-NT-003`,
      title: 'Wire into existing infrastructure',
      layer: 0,
      dependencies: [`${parentId}-NT-002`]
    }
  ];

  return tasks;
}

// — Startup

async function main() {

```

```
const transport = new StdioServerTransport();
await server.connect(transport);
console.error('SuperBridge MCP Server started');
}

main().catch(err => {
  console.error('FATAL:', err);
  process.exit(1);
});
```

G. Implementation Gate – `implementation-gate.js`

The full source code for `implementation-gate.js` is provided in [TDD_INTEGRATION.md §10](#). It includes the RED, GREEN, and REFACTOR phase gates, anti-bluff verification, and deterministic consistency enforcement. Reference it from there as the canonical implementation.

H. Task Granulator – `granulator.ts`

File: `src/granulator.ts`

```
#!/usr/bin/env node
```

```
/**
```

```
 * Task Granulator — Nano-Task Decomposition Engine
```

```
 *
```

```
 * Decomposes a feature specification into a tree of decoupled nano-tasks
```

```
 * suitable for autonomous implementation by Helix LLM (§11.4.58 PWU pipeline).
```

```
 *
```

```
 * Constitutional anchors: §11.4.224 (TDD for ALL work), §11.4.58 (parallel PWU)
```

```
 */
```

```
interface GranulatorConfig {
```

```
    maxLayerDepth: number;
```

```
    maxTasksPerLayer: number;
```

```
    minTaskGranularity: number; // minimum lines of code per nano-task
```

```
    language: string;
```

```
    parentId: string;
```

```
}
```

```
interface NanoTaskSpec {
```

```
    id: string;
```

```
    parent_feature_id: string;
```

```
    title: string;
```

```
    description: string;
```

```
    language: string;
```

```
    test_file: string;
```

```
    impl_file: string;
```

```
    layer: number;
```

```
    dependencies: string[];
```

```
    acceptance_criteria: { description: string; evidence_class: string }[];
```

```
    estimated_tokens: number;
```

```
}
```

```
class TaskGranulator {
```

```
    private config: GranulatorConfig;
```

```
    private taskCounter: Map<number, number> = new Map(); // layer → counter
```

```
    constructor(config: GranulatorConfig) {
```

```
        this.config = config;
```

```
    }
```

```
/**
```

```
 * Granulate a feature specification into nano-tasks.
```

```
 *
```

```
 * @param spec - The feature specification text from SpecKit's specify phase
```

```

* @returns Array of nano-task specifications with dependency ordering
*/
granulate(spec: string): NanoTaskSpec[] {
  const tasks: NanoTaskSpec[] = [];

  // Phase 1: Parse the spec into logical work units
  const workUnits = this.parseSpec(spec);

  // Phase 2: Decompose each work unit into nano-tasks
  for (const unit of workUnits) {
    const unitTasks = this.decomposeUnit(unit, 0);
    tasks.push(...unitTasks);
  }

  // Phase 3: Compute dependencies between nano-tasks
  this.computeDependencies(tasks);

  // Phase 4: Validate the decomposition
  this.validateDecomposition(tasks);

  // Phase 5: Sort by layer and dependency order
  tasks.sort((a, b) => {
    if (a.layer !== b.layer) return a.layer - b.layer;
    return a.id.localeCompare(b.id);
  });

  return tasks;
}

private parseSpec(spec: string): WorkUnit[] {
  // In production: parse the SpecKit spec (user stories, acceptance criteria, technical notes)
  // into a structured set of work units.
  //
  // Each work unit is a self-contained piece of work that can be decomposed
  // into nano-tasks. Work units are ordered by dependency (earlier units
  // produce types/interfaces/foundations that later units consume).

  // Stub implementation: split on double-newline paragraphs
  const paragraphs = spec.split(/\n\n+/).filter(p => p.trim().length > 0);
  const units: WorkUnit[] = [];

  for (let i = 0; i < paragraphs.length; i++) {
    units.push({
      id: i,
      description: paragraphs[i].trim(),
      dependencies: i > 0 ? [i - 1] : [],
    });
  }
}

```

```

        estimatedComplexity: paragraphs[i].length / 100
    });
}

return units;
}

private decomposeUnit(unit: WorkUnit, parentLayer: number): NanoTaskSpec[] {
    const tasks: NanoTaskSpec[] = [];
    const complexity = unit.estimatedComplexity;

    // How many nano-tasks does this unit need?
    const taskCount = Math.max(1, Math.min(
        Math.ceil(complexity / this.config.minTaskGranularity),
        this.config.maxTasksPerLayer
    ));

    const layer = parentLayer;

    for (let i = 0; i < taskCount; i++) {
        if (layer > this.config.maxLayerDepth) break;

        const seq = this.nextSeq(layer);
        const id = `${this.config.parentId}-NT-${String(layer).padStart(3, '0')}${String(seq).padStart(3, '0')}`;

        tasks.push({
            id,
            parent_feature_id: this.config.parentId,
            title: this.deriveTitle(unit, i, taskCount),
            description: this.deriveDescription(unit, i, taskCount),
            language: this.config.language,
            test_file: this.deriveTestPath(id),
            impl_file: this.deriveImplPath(id),
            layer,
            dependencies: [], // computed in phase 3
            acceptance_criteria: this.deriveAcceptanceCriteria(unit, i, taskCount),
            estimated_tokens: Math.ceil(complexity * 200 / taskCount)
        });
    }

    // If the task is still too complex, recurse into deeper layers
    if (layer < this.config.maxLayerDepth && complexity > this.config.minTaskGranularity * 2) {
        for (let i = 0; i < Math.min(taskCount, 2); i++) {
            const subUnit: WorkUnit = {
                id: unit.id * 100 + i,

```

```

        description: `Sub-task of ${unit.description}: part ${i + 1}`,
        dependencies: [],
        estimatedComplexity: complexity / taskCount
    };

    const subTasks = this.decomposeUnit(subUnit, layer + 1);

    // Wire dependencies: sub-tasks depend on their parent layer's task
    if (tasks.length > 0) {
        const parentTaskId = `${this.config.parentId}-NT-${String(layer).padStart(3, '0')}${String(i + 1).padStart(3, '0')}`;
        for (const subTask of subTasks) {
            if (!subTask.dependencies.includes(parentTaskId)) {
                subTask.dependencies.push(parentTaskId);
            }
        }
    }

    tasks.push(...subTasks);
}

return tasks;
}

private computeDependencies(tasks: NanoTaskSpec[]): void {
    const byLayer = new Map<number, NanoTaskSpec[]>();
    for (const task of tasks) {
        if (!byLayer.has(task.layer)) byLayer.set(task.layer, []);
        byLayer.get(task.layer)!.push(task);
    }

    // Layer N tasks depend on layer N-1 tasks (foundational layers first)
    for (let layer = 1; layer <= this.config.maxLayerDepth; layer++) {
        const currentLayer = byLayer.get(layer) || [];
        const prevLayer = byLayer.get(layer - 1) || [];

        for (const task of currentLayer) {
            // Each task in layer N depends on at least one task in layer N-1
            if (prevLayer.length > 0) {
                const depIndex = Math.min(
                    currentLayer.indexOf(task),
                    prevLayer.length - 1
                );
                task.dependencies.push(prevLayer[depIndex].id);
            }
        }
    }
}

```

```

}

// Same-layer sequential tasks depend on the previous task
for (const [, layerTasks] of byLayer) {
  for (let i = 1; i < layerTasks.length; i++) {
    if (!layerTasks[i].dependencies.includes(layerTasks[i - 1].id)) {
      layerTasks[i].dependencies.push(layerTasks[i - 1].id);
    }
  }
}
}

private validateDecomposition(tasks: NanoTaskSpec[]): void {
  const ids = new Set(tasks.map(t => t.id));

  // Check all dependency IDs exist
  for (const task of tasks) {
    for (const dep of task.dependencies) {
      if (!ids.has(dep)) {
        throw new Error(`VALIDATION_FAILED: Task ${task.id} depends on ${dep} which does not exist`);
      }
    }
  }

  // Check for self-dependency
  for (const task of tasks) {
    if (task.dependencies.includes(task.id)) {
      throw new Error(`VALIDATION_FAILED: Task ${task.id} depends on itself`);
    }
  }

  // Check description meets §11.4.171 minimum (50 chars)
  for (const task of tasks) {
    if (task.description.length < 50) {
      throw new Error(`VALIDATION_FAILED: Task ${task.id} description too short (${task.description.length} < 50 chars per §11.4.171)`);
    }
  }
}

private nextSeq(layer: number): number {
  const current = this.taskCounter.get(layer) || 0;
  this.taskCounter.set(layer, current + 1);
  return current + 1;
}

```

```

private deriveTitle(unit: WorkUnit, index: number, total: number): string {
  const prefix = unit.description.split('.')[0] || unit.description;
  return `${prefix.substring(0, 50)} (part ${index + 1}/${total})`;
}

private deriveDescription(unit: WorkUnit, index: number, total: number): string {
  return `Part ${index + 1} of ${total} for: ${unit.description}. ` +
    `This nano-task is a decoupled, independently testable unit of work ` +
    `that contributes to the parent feature ${this.config.parentId}. ` +
    `It must be implemented following TDD: write the test first, observe it fail (RED), ` +
    `then implement until the test passes (GREEN). ` +
    `The task is small enough that a modest local LLM can execute it correctly.`;
}

private deriveTestPath(id: string): string {
  return `nano_tasks/${this.config.parentId}/${id}_test.${this.fileExtension()}`;
}

private deriveImplPath(id: string): string {
  return `nano_tasks/${this.config.parentId}/${id}.${this.fileExtension()}`;
}

private fileExtension(): string {
  switch (this.config.language) {
    case 'go': return 'go';
    case 'typescript': return 'ts';
    case 'javascript': return 'js';
    case 'bash': return 'sh';
    default: return 'ts';
  }
}

private deriveAcceptanceCriteria(unit: WorkUnit, index: number, total: number): { description:
  string; evidence_class: string }[] {
  const criteria: { description: string; evidence_class: string }[] = [];

  criteria.push({
    description: `Nano-task ${index + 1}/${total} implements its specified sub-function correctly`,
    evidence_class: 'source'
  });

  if (total <= 3) {
    criteria.push({
      description: `The implementation passes all assertions in the paired test file`,
      evidence_class: 'runtime'
    });
  }
}

```

```

    }

    return criteria;
}
}

// — Types

interface WorkUnit {
    id: number;
    description: string;
    dependencies: number[];
    estimatedComplexity: number;
}

// — CLI

async function main() {
    const args = process.argv.slice(2);
    if (args.length < 2) {
        console.error('Usage: granulator <parent_id> <spec_file> [--language=go] [--max-depth=3]');
        process.exit(1);
    }

    const parentId = args[0];
    const specFile = args[1];
    const fs = await import('fs');

    if (!fs.existsSync(specFile)) {
        console.error(`Spec file not found: ${specFile}`);
        process.exit(1);
    }

    const spec = fs.readFileSync(specFile, 'utf-8');

    const config: GranulatorConfig = {
        maxLayerDepth: 3,
        maxTasksPerLayer: 10,
        minTaskGranularity: 5,
        language: 'go',
        parentId
    };

```

```

// Parse optional flags
for (const arg of args.slice(2)) {
  if (arg.startsWith('--language=')) config.language = arg.split('=')[1];
  if (arg.startsWith('--max-depth=')) config.maxLayerDepth = parseInt(arg.split('=')[1], 10);
}

const granulator = new TaskGranulator(config);
const tasks = granulator.granulate(spec);

console.log(JSON.stringify({
  parent_id: parentId,
  task_count: tasks.length,
  max_layer: Math.max(...tasks.map(t => t.layer)),
  tasks
}, null, 2));
}

main().catch(err => {
  console.error('FATAL:', err.message);
  process.exit(1);
});

```

I. DAG Engine — dag.go

File: pkg/dag/dag.go

```

package dag

import (
    "errors"
    "fmt"
    "sort"
    "sync"
)

// Node represents a single nano-task in the dependency DAG.
type Node struct {
    ID      string `json:"id"`
    Title   string `json:"title"`
    Layer   int    `json:"layer"`
    Dependencies []string `json:"dependencies"`
    Status  Status `json:"status"`
}

// Status is the execution status of a DAG node.
type Status string

const (
    StatusPending Status = "pending"
    StatusReady   Status = "ready"
    StatusRunning Status = "running"
    StatusDone    Status = "done"
    StatusFailed  Status = "failed"
    StatusBlocked Status = "blocked"
)

// DAG is a directed acyclic graph of nano-tasks.
type DAG struct {
    mu      sync.RWMutex
    Nodes   map[string]*Node `json:"nodes"`
    Adjacency map[string][]string `json:"adjacency" // node → dependents`
    InDegree map[string]int    `json:"in_degree"`
    ReadyQueue []string `json:"ready_queue"`
}

// New creates a new DAG from a set of nodes.
func New(nodes []Node) (*DAG, error) {
    d := &DAG{
        Nodes:   make(map[string]*Node),
        Adjacency: make(map[string][]string),
        InDegree: make(map[string]int),
    }

```

```

}

for i := range nodes {
    n := nodes[i]
    n.Status = StatusPending
    d.Nodes[n.ID] = &n
    d.InDegree[n.ID] = 0
    if d.Adjacency[n.ID] == nil {
        d.Adjacency[n.ID] = []string{}
    }
}

// Build adjacency list and in-degree counts
for _, n := range d.Nodes {
    for _, depID := range n.Dependencies {
        if _, ok := d.Nodes[depID]; !ok {
            return nil, fmt.Errorf("UNRESOLVED_DEPENDENCY: node %s depends on %s which does not exist", n.ID, depID)
        }
        d.Adjacency[depID] = append(d.Adjacency[depID], n.ID)
        d.InDegree[n.ID]++
    }
}

// Cycle detection
if err := d.detectCycle(); err != nil {
    return nil, err
}

// Compute initial ready queue (in-degree == 0)
d.rebuildReadyQueue()

return d, nil
}

// detectCycle uses DFS to detect cycles in O(V+E).
func (d *DAG) detectCycle() error {
    const (
        white = 0 // unvisited
        gray  = 1 // in current DFS path
        black = 2 // fully explored
    )

    color := make(map[string]int)
    for id := range d.Nodes {
        color[id] = white
    }
}

```

```

}

var dfs func(id string) error
dfs = func(id string) error {
    color[id] = gray
    for _, neighbor := range d.Adjacency[id] {
        switch color[neighbor] {
        case gray:
            return fmt.Errorf("CYCLE_DETECTED: back-edge from %s to %s", id, neighbor)
        case white:
            if err := dfs(neighbor); err != nil {
                return err
            }
        }
    }
    color[id] = black
    return nil
}

for id := range d.Nodes {
    if color[id] == white {
        if err := dfs(id); err != nil {
            return err
        }
    }
}
return nil
}

// TopologicalSort returns nodes in topological order (Kahn's algorithm).
func (d *DAG) TopologicalSort() ([]string, error) {
    d.mu.RLock()
    defer d.mu.RUnlock()

    inDegree := make(map[string]int)
    for id, deg := range d.InDegree {
        inDegree[id] = deg
    }

    queue := []string{}
    for id, deg := range inDegree {
        if deg == 0 {
            queue = append(queue, id)
        }
    }

    sort.Strings(queue) // Deterministic order

```

```

result := []string{}
for len(queue) > 0 {
    node := queue[0]
    queue = queue[1:]
    result = append(result, node)

    for _, neighbor := range d.Adjacency[node] {
        inDegree[neighbor]--
        if inDegree[neighbor] == 0 {
            queue = append(queue, neighbor)
        }
    }
}

if len(result) != len(d.Nodes) {
    return nil, fmt.Errorf("TOPOLOGICAL_SORT_FAILED: only %d/%d nodes sortable — cycle or missing node",
        len(result), len(d.Nodes))
}

return result, nil
}

// GetReady returns the current set of nodes ready for execution (all deps satisfied).
func (d *DAG) GetReady() []string {
    d.mu.RLock()
    defer d.mu.RUnlock()

    ready := make([]string, len(d.ReadyQueue))
    copy(ready, d.ReadyQueue)
    return ready
}

// MarkDone marks a node as complete and updates dependent nodes' readiness.
func (d *DAG) MarkDone(id string) error {
    d.mu.Lock()
    defer d.mu.Unlock()

    node, ok := d.Nodes[id]
    if !ok {
        return fmt.Errorf("NODE_NOT_FOUND: %s", id)
    }

    if node.Status == StatusDone {
        return nil // Idempotent
    }

```

```

}

if node.Status != StatusRunning && node.Status != StatusReady {
    return fmt.Errorf("INVALID_TRANSITION: cannot mark %s as done from status %s", id,
        node.Status)
}

node.Status = StatusDone

// Decrement in-degree of dependents
for _, dependent := range d.Adjacency[id] {
    d.InDegree[dependent]--
}

d.rebuildReadyQueue()
return nil
}

// MarkFailed marks a node as failed, blocking its dependents.
func (d *DAG) MarkFailed(id string, reason string) error {
    d.mu.Lock()
    defer d.mu.Unlock()

    node, ok := d.Nodes[id]
    if !ok {
        return fmt.Errorf("NODE_NOT_FOUND: %s", id)
    }

    node.Status = StatusFailed

    // Cascade block to all transitive dependents
    for _, dependent := range d.Adjacency[id] {
        d.cascadeBlock(dependent)
    }

    d.rebuildReadyQueue()
    return nil
}

func (d *DAG) cascadeBlock(id string) {
    node := d.Nodes[id]
    if node.Status == StatusDone || node.Status == StatusFailed {
        return
    }
    node.Status = StatusBlocked
    for _, dependent := range d.Adjacency[id] {

```

```

        d.cascadeBlock(dependent)
    }
}

// MarkRunning marks a node as currently executing.
func (d *DAG) MarkRunning(id string) error {
    d.mu.Lock()
    defer d.mu.Unlock()

    node, ok := d.Nodes[id]
    if !ok {
        return fmt.Errorf("NODE_NOT_FOUND: %s", id)
    }

    if node.Status != StatusReady {
        return fmt.Errorf("INVALID_TRANSITION: cannot run %s from status %s (must be ready)", id,
            node.Status)
    }

    node.Status = StatusRunning
    d.rebuildReadyQueue()
    return nil
}

// IsComplete returns true when all nodes are Done or Failed.
func (d *DAG) IsComplete() bool {
    d.mu.RLock()
    defer d.mu.RUnlock()

    for _, node := range d.Nodes {
        if node.Status != StatusDone && node.Status != StatusFailed {
            return false
        }
    }
    return true
}

// Progress returns (done, total, failed) counts.
func (d *DAG) Progress() (done, total, failed int) {
    d.mu.RLock()
    defer d.mu.RUnlock()

    total = len(d.Nodes)
    for _, node := range d.Nodes {
        switch node.Status {
            case StatusDone:

```

```

        done++
        case StatusFailed:
            failed++
    }
}
return
}

// Layers returns nodes grouped by their layer (0-indexed).
func (d *DAG) Layers() map[int][]string {
    d.mu.RLock()
    defer d.mu.RUnlock()

    layers := make(map[int][]string)
    for id, node := range d.Nodes {
        layers[node.Layer] = append(layers[node.Layer], id)
    }
    for layer := range layers {
        sort.Strings(layers[layer])
    }
    return layers
}

// ValidateSchema validates that every node conforms to the nano-task JSON schema.
func (d *DAG) ValidateSchema() []error {
    d.mu.RLock()
    defer d.mu.RUnlock()

    var errs []error
    for id, node := range d.Nodes {
        if id == "" {
            errs = append(errs, errors.New("node has empty ID"))
        }
        if node.Title == "" {
            errs = append(errs, fmt.Errorf("node %s has empty title", id))
        }
        if node.Layer < 0 {
            errs = append(errs, fmt.Errorf("node %s has negative layer %d", id, node.Layer))
        }
    }
    return errs
}

// rebuildReadyQueue recomputes nodes with in-degree == 0 and status == Pending.
// Caller must hold d.mu.
func (d *DAG) rebuildReadyQueue() {

```

```
d.ReadyQueue = nil
for id, node := range d.Nodes {
    if d.InDegree[id] == 0 && node.Status == StatusPending {
        d.ReadyQueue = append(d.ReadyQueue, id)
        node.Status = StatusReady
    }
}
sort.Strings(d.ReadyQueue)
}
```

J. Constitution-Aware Commands – Bash Scripts

File: `scripts/tdd_cycle.sh`

```
#!/usr/bin/env bash

set -euo pipefail

# tdd_cycle.sh — Run a complete TDD cycle for a nano-task
# §11.4.224: TDD for ALL work. §11.4.115: RED on broken artifact.
#
# Usage: bash scripts/tdd_cycle.sh <item_id> <test_file> <impl_file> [--language=go|ts|bash]
# Environment:
# RED_MODE=1 (default) — runs RED phase first
# TDD_ITERATIONS=3 — deterministic consistency iterations
# TDD_VERDICT_DIR — where to write verdict JSONs
# TDD_EVIDENCE_DIR — where to write captured evidence

ITEM_ID="${1:?Usage: tdd_cycle.sh <item_id> <test_file> <impl_file>}"
TEST_FILE="${2:?}"
IMPL_FILE="${3:?}"
LANG="${4:-go}"
RED_MODE="${RED_MODE:-1}"
ITERATIONS="${TDD_ITERATIONS:-3}"
VERDICT_DIR="${TDD_VERDICT_DIR:-qa-results/tdd_verdicts}"
EVIDENCE_DIR="${TDD_EVIDENCE_DIR:-qa-results/tdd_evidence}"

# Strip --language= prefix if present
LANG="${LANG#--language=}"

mkdir -p "$VERDICT_DIR" "$EVIDENCE_DIR/$ITEM_ID"

# ——— Helpers

sha256_of() {
    # Compute fingerprint of item + files
    local content="${ITEM_ID}"
    for f in "$@"; do
        if [[ -f "$f" ]]; then
            content+=$(<"$f")
        fi
    done
    echo -n "$content" | sha256sum | cut -d' ' -f1
}

now_iso() {
    date -u +"%Y-%m-%dT%H:%M:%SZ"
}
```

```

run_test() {
    local mode="$1"
    local lang="$2"
    local test_file="$3"

    case "$lang" in
        go)
            RED_MODE="$mode" go test -v -count=1 "$test_file" 2>&1 || true
            ;;
        ts|typescript)
            RED_MODE="$mode" npx vitest run "$test_file" --reporter=verbose 2>&1 || true
            ;;
        bash|sh)
            RED_MODE="$mode" bash -euo pipefail "$test_file" 2>&1 || true
            ;;
        *)
            echo "UNSUPPORTED_LANGUAGE: $lang" >&2
            return 99
            ;;
    esac
}

```

— RED Phase

```

run_red() {
    echo "=== RED PHASE: $ITEM_ID ==="

    if [[ ! -f "$TEST_FILE" ]]; then
        echo "RED_FAILED: Test file $TEST_FILE does not exist. Write the test first."
        return 1
    fi

    if [[ -f "$IMPL_FILE" ]]; then
        echo "RED_FAILED: Implementation file $IMPL_FILE already exists. Delete it and write the test first."
        return 1
    fi

    local red_output
    red_output=$(run_test "1" "$LANG" "$TEST_FILE")
    local red_rc=$?

    if [[ $red_rc -eq 0 ]]; then
        echo "RED_FAILED: Test passed in RED mode — tautology. The test does not catch the gap."
        return 1
    fi
}

```

fi

Deterministic consistency: run N iterations

local exit_codes=()

for ((i=1; i<=ITERATIONS; i++)); **do**

set +e

 run_test "1" "\$LANG" "\$TEST_FILE" > /dev/null 2>&1

 exit_codes+=(\$?)

set -e

done

local unique_codes

unique_codes=\$(**printf** '%s\n' "\${exit_codes[@]}" | **sort** -u)

if [[\$(**echo** "\$unique_codes" | **wc** -l) -ne 1]]; **then**

echo "RED_FAILED: Nondeterministic: exit codes across \$ITERATIONS iterations: \$
 \${exit_codes[*]}"

return 1

fi

Write evidence

local evidence_file="\$EVIDENCE_DIR/\$ITEM_ID/RED_\${ITEM_ID}_stdout.txt"

echo "\$red_output" > "\$evidence_file"

Write RED verdict

local fingerprint

fingerprint=\$(sha256_of "\$TEST_FILE")

local verdict_file="\$VERDICT_DIR/RED_\${ITEM_ID}.json"

cat > "\$verdict_file" <<JSONEOF

```
{
  "item_id": "$ITEM_ID",
  "guard_identity": "tdd_gate_${ITEM_ID}",
  "polarity": "RED",
  "exit_code": $red_rc,
  "target_fingerprint": "$fingerprint",
  "iterations": $ITERATIONS,
  "evidence_class": "source",
  "evidence_files": ["$evidence_file"],
  "precondition_provenance": "observed",
  "timestamp": "$(now_iso)"
}
```

JSONEOF

echo "RED_PASS: Verdict written to \$verdict_file"

echo "RED_PASS: Evidence written to \$evidence_file"

return 0

```
}
```

```
# ——— GREEN Phase
```

```
run_green() {  
    echo "=== GREEN PHASE: $ITEM_ID ==="  
  
    local red_verdict="$VERDICT_DIR/RED_${ITEM_ID}.json"  
    if [[ ! -f "$red_verdict" ]]; then  
        echo "GREEN_REJECTED: No prior RED verdict at $red_verdict. Run RED phase first."  
        return 1  
    fi  
  
    local red_fingerprint  
    red_fingerprint=$(jq -r '.target_fingerprint' "$red_verdict")  
  
    if [[ ! -f "$IMPL_FILE" ]]; then  
        echo "GREEN_REJECTED: Implementation file $IMPL_FILE does not exist."  
        return 1  
    fi  
  
    local green_output  
    green_output=$(run_test "0" "$LANG" "$TEST_FILE")  
    local green_rc=$?  
  
    if [[ $green_rc -ne 0 ]]; then  
        echo "GREEN_FAILED: Test did not pass (exit $green_rc)."  
        return 1  
    fi  
  
    local green_fingerprint  
    green_fingerprint=$(sha256_of "$TEST_FILE" "$IMPL_FILE")  
  
    if [[ "$green_fingerprint" == "$red_fingerprint" ]]; then  
        echo "GREEN_FAILED: Identical fingerprints — the artifact has not changed since RED. Fix not  
            deployed."  
        return 1  
    fi  
  
    # Deterministic consistency  
    local exit_codes=()  
    for ((i=1; i<=ITERATIONS; i++)); do  
        set +e  
        run_test "0" "$LANG" "$TEST_FILE" > /dev/null 2>&1  
        exit_codes+=($?)  
    done  
}
```

```

set -e
done

local unique_codes
unique_codes=$(printf '%s\n' "${exit_codes[@]}" | sort -u)
if [[ "$unique_codes" != "0" ]]; then
    echo "GREEN_FAILED: Nondeterministic: not all iterations passed (exit codes: ${exit_codes[*]})"
    return 1
fi

# Write evidence
local evidence_file="$EVIDENCE_DIR/$ITEM_ID/GREEN_${ITEM_ID}_stdout.txt"
echo "$green_output" > "$evidence_file"

# Write GREEN verdict
local verdict_file="$VERDICT_DIR/GREEN_${ITEM_ID}.json"
cat > "$verdict_file" <<JSONEOF
{
    "item_id": "$ITEM_ID",
    "guard_identity": "tdd_gate_${ITEM_ID}",
    "polarity": "GREEN",
    "exit_code": 0,
    "target_fingerprint": "$green_fingerprint",
    "iterations": $ITERATIONS,
    "evidence_class": "source",
    "evidence_files": ["$evidence_file"],
    "prior_red_verdict": "$red_verdict",
    "timestamp": "$(now_iso)"
}
JSONEOF

echo "GREEN_PASS: Verdict written to $verdict_file"
return 0
}

# — Main

if [[ "$RED_MODE" == "1" ]]; then
    run_red || exit 1
fi

run_green || exit 1

echo ""

```

```
echo "===== "  
echo "TDD CYCLE COMPLETE: $ITEM_ID"  
echo " RED verdict: $VERDICT_DIR/RED_${ITEM_ID}.json"  
echo " GREEN verdict: $VERDICT_DIR/GREEN_${ITEM_ID}.json"  
echo " Evidence: $EVIDENCE_DIR/$ITEM_ID/"  
echo "===== "
```

K. Pack Distribution Script – `pack.sh`

File: `scripts/pack.sh`

```
#!/usr/bin/env bash
set -euo pipefail

# pack.sh — Package the Speckit-Superpowers bridge for distribution
# Produces a self-contained tarball with all source, configs, and docs.
#
# §11.4.65: Universal Markdown export — all docs get .html+.pdf twins.
# §11.4.73: Project-styling mandate.

VERSION="${1:?Usage: pack.sh <version> (e.g., 1.0.0)}"
OUTPUT_DIR="dist"
PACK_NAME="speckit-superpowers-bridge-${VERSION}"
PACK_DIR="${OUTPUT_DIR}/${PACK_NAME}"

echo "Packing Speckit-Superpowers Bridge v${VERSION}..."

# Clean previous build
rm -rf "${OUTPUT_DIR:?}/${PACK_NAME:?}"

# Create directory structure
mkdir -p "$PACK_DIR"/{src,cmd,pkg,scripts,schemas,config,docs,certs,containers}

# Copy source files
cp src/index.ts "$PACK_DIR/src/"
cp src/granulator.ts "$PACK_DIR/src/"
cp src/implementation-gate.ts "$PACK_DIR/src/"

# Copy Go source
cp -r cmd/ "$PACK_DIR/"
cp -r pkg/ "$PACK_DIR/"

# Copy scripts
cp scripts/tdd_cycle.sh "$PACK_DIR/scripts/"
cp scripts/pack.sh "$PACK_DIR/scripts/"
chmod +x "$PACK_DIR/scripts/*.sh"

# Copy schemas
cp schemas/nano-task.schema.json "$PACK_DIR/schemas/"

# Copy configs
cp mcp.json "$PACK_DIR/config/"
cp extension.yml "$PACK_DIR/config/"
cp .env.template "$PACK_DIR/config/"

# Copy documentation
```

```

# Generate HTML + PDF twins for all Markdown docs
echo "Generating documentation exports..."
for md_file in TDD_INTEGRATION.md SECURITY.md APPENDIX.md; do
    if [[ -f "docs/$md_file" ]]; then
        cp "docs/$md_file" "$PACK_DIR/docs/"

        # Generate HTML
        if command -v pandoc &>/dev/null; then
            pandoc "docs/$md_file" \
                --standalone \
                --from gfm \
                --to html5 \
                --metadata title="$(head -1 "docs/$md_file" | sed 's/^# //' )" \
                --output "$PACK_DIR/docs/${md_file%.md}.html" \
                2>/dev/null || echo "WARNING: pandoc failed for $md_file"
        fi

        # Generate PDF
        if command -v pandoc &>/dev/null && command -v weasyprint &>/dev/null; then
            pandoc "docs/$md_file" \
                --from gfm \
                --to html5 \
                --metadata title="$(head -1 "docs/$md_file" | sed 's/^# //' )" \
                --output /tmp/${$$_speckit_doc}.html 2>/dev/null && \
            weasyprint /tmp/${$$_speckit_doc}.html "$PACK_DIR/docs/${md_file%.md}.pdf" 2>/dev/null && \
            rm -f /tmp/${$$_speckit_doc}.html || \
            echo "WARNING: PDF generation failed for $md_file"
        fi
    fi
done

# Generate DOCX where possible
for md_file in TDD_INTEGRATION.md SECURITY.md APPENDIX.md; do
    if [[ -f "docs/$md_file" ]] && command -v pandoc &>/dev/null; then
        pandoc "docs/$md_file" \
            --from gfm \
            --to docx \
            --output "$PACK_DIR/docs/${md_file%.md}.docx" 2>/dev/null || \
            echo "WARNING: DOCX generation failed for $md_file"
    fi
done

# Copy README
if [[ -f README.md ]]; then
    cp README.md "$PACK_DIR/"

```

fi

Create tarball

echo "Creating tarball..."

tar -czf "\${OUTPUT_DIR}/\${PACK_NAME}.tar.gz" -C "\$OUTPUT_DIR" "\$PACK_NAME"

Create SHA256 checksum

sha256sum "\${OUTPUT_DIR}/\${PACK_NAME}.tar.gz" > "\${OUTPUT_DIR}/\${PACK_NAME}.tar.gz.sha256"

echo ""

echo "Pack complete:"

echo " Tarball: \${OUTPUT_DIR}/\${PACK_NAME}.tar.gz"

echo " SHA256: \$(cat "\${OUTPUT_DIR}/\${PACK_NAME}.tar.gz.sha256")"

echo " Size: \$(du -h "\${OUTPUT_DIR}/\${PACK_NAME}.tar.gz" | cut -f1)"

L. Environment Template – `.env.template`

```

# =====

# SpecKit-Superpowers Bridge — Environment Template
# =====

# Copy this file to .env and fill in the values.
# .env is git-ignored per §11.4.30. NEVER commit real credentials.
# =====

# ——— Helix LLM Gateway


---


HELIX_GATEWAY_URL=https://localhost:3101
HELIX_GATEWAY_TLS_CERT=/path/to/certs/gateway-client.crt
HELIX_GATEWAY_TLS_KEY=/path/to/certs/gateway-client.key
HELIX_GATEWAY_CA_CERT=/path/to/certs/ca.crt

# ——— llama.cpp RPC


---


LLAMA_RPC_HOST=localhost
LLAMA_RPC_PORT=8080
LLAMA_RPC_TLS_ENABLED=false
LLAMA_MODEL_PATH=/models/llama-3-8b-instruct.Q4_K_M.gguf
LLAMA_MODEL_SHA256= # Verified model hash from vendor — fill in

# ——— TDD Enforcement


---


TDD_VERDICT_DIR=qa-results/tdd_verdicts
TDD_EVIDENCE_DIR=qa-results/tdd_evidence
TDD_ITERATIONS=3
TDD_COVERAGE_THRESHOLD=0.85
TDD_MAX_CYCLES=5
TDD_SANDBOX_ENABLED=true
TDD_TIMEOUT_SECONDS=300

# ——— Container Runtime §11.4.161


---


CONTAINER_RUNTIME=podman
CONTAINER_ROOTLESS=true
NANO_TASK_SANDBOX_IMAGE=nano-task-runner:latest
NANO_TASK_MEM_LIMIT=256m
NANO_TASK_CPU_LIMIT=0.5

# ——— External Services


---



```

Cloud LLM provider API keys (only used as last-resort fallback per §11.4.196)

CLOUD_LLM_PROVIDER= *# e.g., openai, anthropic — optional*

CLOUD_LLM_API_KEY= *# NEVER commit*

— Project Identity §11.4.151

HELIX_RELEASE_PREFIX=speckit-bridge

— ClickUp Integration §11.4.148(D5)

CLICKUP_API_TOKEN= *# NEVER commit*

CLICKUP_LIST_ID= *# Workable items list*

CLICKUP_WORKSPACE_ID= *# Workspace ID*

CLICKUP_ASSIGNEE= *# Default assignee username*

— GitHub/GitLab CLIs

GITHUB_TOKEN= *# NEVER commit*

GITLAB_TOKEN= *# NEVER commit*

M. Container Compose Files

File: containers/docker-compose.yml

```
# docker-compose.yml — SpecKit-Superpowers Bridge Services
# All services run rootless per §11.4.161.
# Provisioned via vasic-digital/containers submodule §11.4.76.
```

```
version: "3.9"
```

```
services:
```

```
# ——— Helix LLM Gateway
```

```
helix-gateway:
```

```
image: "helix-gateway:latest"
```

```
build:
```

```
context: .
```

```
dockerfile: containers/Dockerfile.gateway
```

```
ports:
```

```
- "3101:3101"
```

```
volumes:
```

```
- ./config/gateway.json:/app/config/gateway.json:ro
```

```
- ./certs:/app/certs:ro
```

```
- ./models:/models:ro
```

```
environment:
```

```
- GATEWAY_CONFIG=/app/config/gateway.json
```

```
restart: unless-stopped
```

```
security_opt:
```

```
- no-new-privileges:true
```

```
cap_drop:
```

```
- ALL
```

```
user: "1000:1000"
```

```
mem_limit: "512m"
```

```
cpus: "1.0"
```

```
# ——— llama.cpp Inference Server
```

```
llama-server:
```

```
image: "llama-cpp-server:latest"
```

```
build:
```

```
context: .
```

```
dockerfile: containers/Dockerfile.llama
```

```
ports:
```

```
- "8080:8080"
```

```
volumes:
```

```
- ./models:/models:ro
```

```
environment:
```

```
- LLAMA_MODEL=/models/llama-3-8b-instruct.Q4_K_M.gguf
```

```
- LLAMA_HOST=0.0.0.0
```

```
- LLAMA_PORT=8080
```

```
- LLAMA_N_GPU_LAYERS=35
```

```
deploy:
  resources:
    reservations:
      devices:
        - driver: nvidia
          count: 1
          capabilities: [gpu]
  restart: unless-stopped
security_opt:
  - no-new-privileges:true
cap_drop:
  - ALL
user: "1000:1000"
mem_limit: "24g"

# ——— Audit Log Aggregator ———
audit-logger:
  image: "audit-logger:latest"
  build:
    context: .
    dockerfile: containers/Dockerfile.audit
  volumes:
    - ./qa-results/audit:/var/log/audit
  environment:
    - AUDIT_LOG_DIR=/var/log/audit
    - AUDIT_RETENTION_DAYS=90
  restart: unless-stopped
  security_opt:
    - no-new-privileges:true
  cap_drop:
    - ALL
  user: "1000:1000"
  mem_limit: "128m"
```

File: `containers/Dockerfile.gateway`

Dockerfile.gateway — Helix LLM Gateway build

FROM golang:1.22-alpine **AS** builder

WORKDIR /build

COPY go.mod go.sum ./

RUN go mod download

COPY cmd/ ./cmd/

COPY pkg/ ./pkg/

RUN CGO_ENABLED=0 GOOS=linux go build -o /gateway ./cmd/gateway

FROM alpine:3.20

RUN apk add --no-cache ca-certificates tzdata

COPY --from=builder /gateway /usr/local/bin/gateway

USER 1000:1000

EXPOSE 3101

ENTRYPOINT ["/usr/local/bin/gateway"]

CMD ["--config", "/app/config/gateway.json"]

N. Test Artifact Template

File: `templates/nano_task_test_template.go.tmpl`

```

package nano_task

import (
    "os"
    "testing"
)

// TestNanoTask_{{.ItemID}} is the TDD guard for nano-task {{.ItemID}}.
// Generated from speckit-superpowers-bridge template.
// RED_MODE=1: assert the defect/gap is present (no implementation yet).
// RED_MODE=0: assert the implementation is correct (§11.4.115 polarity switch).

func TestNanoTask_{{.ItemID}}(t *testing.T) {
    redMode := os.Getenv("RED_MODE")
    if redMode == "" {
        redMode = "1"
    }

    // Precondition: implementation is absent or broken in RED mode
    // The test must genuinely catch the gap — not be a tautology.

    // Arrange
    input := "{{.TestInput}}"
    expected := "{{.ExpectedOutput}}"

    // Act
    result, err := {{.FunctionName}}(input)

    // Assert — RED mode
    if redMode == "1" {
        if err == nil && result == expected {
            t.Fatalf("RED_TAUTOLOGY: {{.FunctionName}} returned correct result in RED mode. " +
                "The test is not catching the gap. Is the implementation already present?")
        }
        t.Logf("RED_PASS: {{.FunctionName}} correctly reports broken/absent in RED mode (err=%v)",
            err)
        return
    }

    // Assert — GREEN mode
    if err != nil {
        t.Fatalf("GREEN_FAILED: {{.FunctionName}} returned error: %v", err)
    }
    if result != expected {
        t.Fatalf("GREEN_FAILED: {{.FunctionName}}(%q) = %q, want %q", input, result, expected)
    }
}

```

```
t.Logf("GREEN_PASS: {{.FunctionName}}(%q) = %q", input, result)
}
```

1 1 4 1 5 6

0. CI Configuration – Disabled per \$. . .

File: `.github/workflows/constitution-compliance.yml.disabled-local-only`

```

# =====

# DISABLED — per §11.4.156: All CI/CD automation MUST be disabled.
# =====

# This file is deliberately named .disabled-local-only so GitHub Actions
# does NOT execute it. Remote CI is DISABLED; enforcement migrates to
# LOCAL git hooks (§11.4.75) + pre-build verification (§11.4.40).
#
# NEVER rename this to .yaml — that would re-enable remote CI, which is a
# §11.4.156 violation and a release blocker.
# =====

#
# What this workflow DID before being disabled (preserved for reference):
#
# name: Constitution Compliance
# on:
#   push:
#     branches: [main]
#   pull_request:
#     branches: [main]
#
# jobs:
#   pre-build-verification:
#     runs-on: ubuntu-latest
#     steps:
#       - uses: actions/checkout@v4
#       with:
#         submodules: recursive
#       - name: Run pre-build gates
#         run: bash constitution/scripts/verify-all-constitution-rules.sh
#
#   meta-test:
#     runs-on: ubuntu-latest
#     steps:
#       - uses: actions/checkout@v4
#       with:
#         submodules: recursive
#       - name: Run meta-test mutations
#         run: bash scripts/testing/meta_test_false_positive_proof.sh
#
#   propagation-gates:
#     runs-on: ubuntu-latest
#     steps:

```

```
# - uses: actions/checkout@v4
#   with:
#     submodules: recursive
# - name: Verify propagation
#   run: bash constitution/scripts/gates/CM-COVENANT-*.sh
#
# =====

# All CI jobs are now LOCAL-ONLY per §11.4.156 + §11.4.75 Layer 5.
# =====
```

End of Appendix.