

OPERATOR_COMPLAINT_MINING

4

Operator Complaint Mining – Corpus (Claude Code session transcripts)

Field	Value
Revision	2
Created	2026-07-22
Last modified	2026-07-22T21:20:00Z
Status	COMPLETE (single-session mining pass; UNKNOWN register in §10)
Author	BG-SESSION-COMPLAINTS forensic mining agent (T1/main - claude4, Fable)
Part of	BG-QUALITY-ROOTCAUSE (see <code>ROOT_CAUSE_ANALYSIS.md</code> , corpora 1–3)

. What this corpus is and why it outranks the tracker

Corpora 1–3 measured what the **tracker recorded**. This corpus measures what the **operator actually experienced**: the first-touch failure reports typed into interactive agent sessions at the moment of discovery, unmediated by any tracker. Phase 1 established that the operator-driven reopens left **no tracker events at all** while

every recorded reopen event was attributed "AI" — so the transcripts are the primary record of the detection channel the tracker never captured. Where complaints and tracker records disagree, the gap IS the finding.

. Corpus definition (measured — every number carries its command)

- Stores: **2 unique** transcript stores.
 - `/home/milos/.claude-shared/projects/` — every per-alias `~/.claude-<alias>/projects` is a symlink to it (verified: `readlink -f` of all `/home/milos/.claude-*/projects` resolves here for every alias; inode check confirmed identity).
 - `/home/milos/.claude/projects/` — the default-alias store, a REAL separate directory (distinct inode), 52 JSONL transcripts incl. `-mnt-track4-atmosphere-t4`.
- Sizes: shared store 7.9 GB (`du -sh`), 486 project dirs (`ls | wc -l`), 16,478 JSONL at full depth (`find -name '*.jsonl' | wc -l`), of which 9,754 at depth ≤ 2 ; the deeper ones are 6,113 `subagents/*.jsonl` (agent-channel), `wf_*` workflow transcripts (SDK channel), and `archive/` real sessions (notably the pre-multi-track atmosphere path `-run-media-milosvasic-DATA4TB-Projects-Android-15`).
- Event schema (verified empirically on real lines, not assumed): one JSON event per line; operator-relevant events have `type=="user"` , `.message.role=="user"` , `.message.content` string (typed prompt) or array (tool results / blocks); `isSidechain==true` marks subagent threads whose "user" messages are **agent-authored**; `entrypoint` distinguishes `cli` (interactive operator) from `sdk-py` / `sdk-cli` (agent-spawned headless workers). `isMeta` marks harness-injected user events.

Operator channel definition used throughout: `type==user` $\wedge$ `role==user` $\wedge$ `entrypoint=="cli"` $\wedge$ `isSidechain!=true` $\wedge$ `isMeta!=true` $\wedge$ path not under `subagents/` $\wedge$ text not a command/hook wrapper (`<command-` , `<bash-` , `<local-command` , Caveat: , [Request interrupted]). Everything else is the agent channel and is excluded from complaint attribution.

. Method

Streaming line-parse (never whole-file reads into context), substring pre-filter before JSON parse, malformed lines counted and reported (never a silent zero). Complaint detection = tiered regex battery over operator-channel prompts, with per-occurrence negation-context checks so standing directive language ("no bluff of any kind", "no false results") does NOT count as a complaint — only affirmative failure reports do. Credential-shaped strings redacted before any excerpt is stored (§11.4.10). Dedup by normalized-text hash (multi-track re-paste of the same complaint counts once per day).

Control needles (§11.4.201(7)(b)): every absence claim in this report carries a known-present needle run through the same pipeline; see §9.

2 0 2 6 0 7 2 2

. The recorded side of the comparison (read-only, measured — —)

Tracker (docs/workable_items.db , opened mode=ro): 619 items;
item_history events: Updated 288, Opened 287, Completed 25, Obsolete 25, **Reopened 25**, Fixed 24, Implemented 2. **All 25 Reopened events are attributed by= 'AI'** (SELECT by, count(*) FROM item_history WHERE event_type LIKE '%eopen%' GROUP BY by) — zero operator-attributed reopen events exist, confirming the Phase-1 finding that the operator detection channel bypasses the tracker entirely. The recorded reopens concentrate on 15 items: ATM-277 ×4 (stale paused-frame on 2nd display), ATM-346/349/328/329 (audio output/focus), ATM-347/353 (player crashes), ATM-350/406 (2nd-display routing), ATM-351 (MPV 4K quality), ATM-352 (subtitles on wrong display), plus 2 design-mockup tasks and 1 DB-tooling item.

Release anchors for clustering (git tags in this checkout, git for-each-ref): dense tag storm 2026-04-18 → 04-29 (1.1.3 → 1.1.5 series), 05-09 → 05-23 (1.1.5-dev-0.0.8 → 14), 05-28 (1.1.6), 05-31 (1.1.7), 06-04 (1.1.8), 06-08 (1.1.9), 07-12 (1.2.0-dev-0.0.1), 07-13 (atmosphere-1.2.1-dev-0.0.4). Known flash events not represented by tags (from the durable session record): 1.2.1-dev-0.0.7 flashed

2026-07-16; 0.0.8 built 2026-07-21. Caveat: tags are gated on validation, so a "release delivered to the operator" is sometimes a flash without a tag; both anchor sets are used in §5.

. Results – headline numbers

Commands: pass-2 extractor `extract_op_prompts.py` (walks both stores, skips `subagents/`, keeps `type==user ^ role==user ^ entrypoint==cli ^ ¬sidechain ^ ¬meta ^ ¬wrapper-prefix`), classifier `classify_op.py` + hand-audited `final_classify.py`. All intermediates under `/tmp/opm/`; the deliverable is `complaints.tsv` beside this file (91 classified rows, redacted).

Measure	Value
JSONL files walked (both stores, full depth)	16,555
Files in operator extraction (subagent transcripts excluded)	9,843
Operator-channel prompts (the entire real operator voice)	2,039
Corpus time coverage (first → last operator prompt)	2026-05-26 → 2026-07-22 (58 days)
Pattern-battery raw matches (operator channel, deduped hash+day)	83
+ ISSUE: / BUG: -prefixed reports the battery missed	8
Classified complaint-candidate rows (deliverable)	91
Class A — verified first-touch defect reports	23
Class B — recurrence meta-complaints ("reopened over and over")	5
Class C — verification-trust complaints ("verified docs are garbage", "flash didn't take")	3
Class D — QA-pressure / directive-context (excluded from complaint totals)	45
Class E — carrier false-positives (pasted covenant / mandate boilerplate)	15
Verification probes (operator questioning claimed verification)	3 unique

Measure	Value
Malformed JSONL lines	0 observed in pass 2; pass 1 (aborted at ~2.3 GB) had recorded 0 to that point

The single most important structural fact: the interactive operator channel is TINY — 2,039 prompts against ~16.5k transcripts. Everything else in the stores is agent-generated. Any mining of this corpus that does not first isolate

`entrypoint=="cli"` and strip the three junk classes (§9) measures the agents talking to themselves, not the operator.

Coverage truncation (load-bearing caveat): nothing before 2026-05-26 survives in either store. The April 1.1.3 → 1.1.5 tag-storm era has ZERO transcript coverage; all absence claims are bounded to the 58-day window.

. Recurrence by subsystem (experienced reopen counts)

Class-A defect reports by primary subsystem (multi-tag texts take the first-priority tag): display_ui 7, video_playback 5, subtitles 4, audio 3, network 1, boot_flash 1, device_stability 1, other(toolkit) 1. Grouping by DEFECT FAMILY rather than tag — the form that maps onto the tracker:

Defect family	Operator-experienced report occasions (dated)	Tracker ids in the family (measured)	Recorded reopen events for the family
Video/sub-titles routed to WRONG display (2nd-display enforcing)	≥5: 2026-05-29 ("subtitles are still being played on 1st display"), 07-16 ×4 (ATM-787/790/792/793/794 sources), 07-17 ("first video clip ... played on primary display instead of 2nd"), 07-21 ×2 (ATM-809 YouTube no picture, ATM-814 Kinopoisk)	≥8: ATM-352, ATM-239, ATM-277, ATM-787, ATM-790, ATM-793, ATM-794, ATM-809, ATM-814	6 (ATM-352 ×2 on 06-09, ATM-277 ×4 on 07-09..14) — all by= 'AI '
Audio output switching	≥3: 05-28 ("5.1 movies ... as stereo"), 07-16 (switch dialog hangs, endless spinner), 07-17 (re-raise)	≥7: ATM-306, ATM-346, ATM-361, ATM-451, ATM-474, ATM-789, ATM-799	2 (ATM-346 ×2 on 06-09); ATM-346 sits in status Reopened since then
4K local playback glitch	3: 05-28 (VLC/MPV compare), 06-09 (FIND-10), 07-16 (ATM-800)	≥2: ATM-351, ATM-800	2 (ATM-351 ×2)
YouTube usable at all (cert/sign-in/playback)	4: 05-29 ("again got same error ... device is not certified"), 07-16 (sign-in as Google Pixel), 07-17, 07-21	≥3 ids	0
Brightness sliders ineffective on connected displays	2: 07-16 (ATM-788), 07-21 (ATM-812 — same defect, NEW id, 5 days later)	2	0

Defect family	Operator-experienced report occasions (dated)	Tracker ids in the family (measured)	Recorded reopen events for the family
exFAT USB mount	≥2: 07-16 titled "REGRESSION: ex-FAT USB flash drive no longer mounts"	≥1 (ATM-797)	0

The experienced-vs-recorded delta. In-window the tracker holds 25 Reopened events (= 18 distinct item-reopens; the 06-09 block double-recorded 7 items), **100% attributed by= 'AI'** . Against that: at least **6 operator-experienced recurrences with ZERO corresponding reopen events** — "All reported issues with Watch here button still exist!", "Subtitles never show on 2nd display ... has not been fixed at all!!!", the 07-17 re-raises (2), ATM-797's own "REGRESSION" title, ATM-812 duplicating ATM-788. This independently reproduces Phase 1's "six operator-driven reopens left no tracker events" from the opposite side of the seam — and the fate of every one of them is visible here: **each entered the tracker as a NEW id** (PC-6 measured live at the intake seam): the 07-16 QA burst minted ATM-786...ATM-800, the 07-21 burst minted ATM-809...ATM-824, including exact same-defect re-mints (ATM-793 $\triangleq$ ATM-352's defect; ATM-812 $\triangleq$ ATM-788; ATM-800 $\triangleq$ ATM-351's family). The `reopens_count` signal of the original ids never moved.

. Post-release clustering

Directly measured, first time for this project. Complaints cluster HARD on flash/QA events, not on tag dates per se:

- **0.0.7 (flashed 2026-07-16):** QA session starts 19:42 with the very first touch broken — "flashed device after flashing is crashing after booting Android! It entered in crashing loop" — followed by **8 ISSUE: reports in 32 minutes (20:50 → 21:22)**; 10 class-A reports on the day. The operator's own 07-19 summary: "Go through all 10+ ISSUES we have reported during manual QA testing of release 0.0.7."

- **0.0.8 (built 07-21, QA same morning):** session opens 09:45 ("We will be reporting issues one by one..."), **first defect at 09:57 — 12 minutes to first failure** (YouTube on 2nd display: sound but no picture), 3 class-A in 32 minutes; 6 QA bug ids minted that day (ATM-809/810/812/814/818/821).
- **Old era, 1.1.6 window (05-26 → 29):** the raw pattern-match rate jumps from 4.88 to 17.19 matches per 100 operator prompts across the 05-28 tag (3.5×), driven by the 05-28/29 manual-test session (VLC 4K glitches, 5.1-as-stereo, seek-freeze, subtitles-on-1st-display, YouTube-cert-again).
- Baseline across the window: 6.35 raw pattern matches per 100 operator prompts (65/1024 atmosphere-family); class-A reports concentrate ~90% within hours of a flash/QA event. **"Broken on first touch" is not a feeling — it is 12 minutes (0.0.8) and 0 minutes (0.0.7 crash-on-boot) measured.**
- Empty windows around the May-09/20/23 and 06-04/08 anchors are corpus truncation or zero-operator-activity days (denominator = 0 shown in the data), NOT evidence of complaint-free releases.

. Complaint→fix→complaint chains (latencies)

1. **Subtitles-on-primary-display** — 05-29 complaint → FIND-5/7 → ATM-352 reopened 06-09 → ATM-239 "In testing" → **07-16 21:02 "Subtitles never show on 2nd display, always on primary one!!! This has not been fixed at all!!!"** (new id ATM-793) → follow-up directives 07-17. ≥3 operator-experienced cycles over 49 days; recorded reopen events for the family: 2, all "AI".
2. **2nd-display routing / watch-here** — ATM-277 recorded reopen-fix cycles ×4 inside ONE week (07-09 → 14, all "AI") → operator first touch 07-16: "All reported issues with Watch here button still exist!" + one-way toggle (ATM-790/794) → 0.0.8 first touch 07-21: ATM-809 (no picture), ATM-814 (Kinopoisk on primary). The feature failed the operator's touch in BOTH releases following its 4th recorded reopen; zero further reopen events.
3. **Audio output switch** — 05-28 "5.1 ... as stereo" → ATM-346 reopened 06-09 (status still Reopened today) → 07-16 20:50 dialog-hang ISSUE → unanswered → **07-17 10:28 operator re-sends the identical ISSUE as**
REMINDER :: on another track (13.6 h re-raise latency) → item ATM-789 minted 07-18 (~1.6 days complaint → item). Family spans ≥7 ids over 51 days.

4. **4K playback glitch** — 05-28 → ATM-351 reopened 06-09 → 07-16 ATM-800 ("mechanism unconfirmed"). 49-day span, 3 occasions, 2 recorded reopens.
5. **Cross-project control (claude-toolkit v1.24.0)** — 07-20 20:07 operator leaves for the night ("You are on your own until the morning ... Having the best possible, most stable build made until the mor[ning]") → **07-20 08:20 (prior morning build) "We cant open any providers alias!"** → 07-22 06:55 "Why helixaget ... cant be even opened!?!?! **Wasnt all this tested before the release!?!"**. Same first-touch shape, different project, no device involved — the failure mode is pipeline-generic, not RK3588-specific.

. Operator-named verification-failure causes vs PC- ..PC-

The operator's lived words map onto the forensic causes — and name two the taxonomy does not cover:

Operator's words (verbatim, dated)	PC mapping
"Wasnt all this tested before the re-release!?!\" (07-22, toolkit)	PC-2 / PC-7 — absence of verification indistinguishable from verification; guards never run where the defect lives
"we have not fixed something properly!\" (07-16, 0.0.7 crash-loop)	PC-3 — source/artifact green read as done
"its flash didn't take (came back on the old Jul-9 build)\" (07-14, quoted agent admission the operator escalates)	PC-3 at the ARTIFACT → RUNTIME seam (the §11.4.200 deploy-verify class) — validated work never reached the device
ATM-819 (07-21, from QA): "TEST-PIPELINE BLUFF: video-to-TV full-auto- mation reports GREEN..."	PC-8 / PC-5 — wrong-layer oracle certifying a user-visible property; the operator filed the pipeline itself as a bug
"All reported issues with Watch here button still exist!\" / "has not been fixed at all!!!!\" (07-16)	PC-6 + PC-3 — recurrence under new ids; re-fix never reached the operator's device in working form
"Documents we have verified are garbage!\" (06-29)	PC-5 family — the verifying instruments/documents themselves false
"All issues reopened multiple times must be fully confirmed, validated and verified during live testing!\" (07-16 05:36 — BEFORE the QA that found them still broken)	PC-1/PC-9 — the operator already distrusts terminal statuses; the same evening proves him right

Causes visible ONLY in this corpus (candidates to add to the taxonomy):

- **NC-1 — External-state drift on an unchanged artifact.** "We again got same error message when tried to use YouTube, device is not certified ... **Yesterday with same flashed image we used to play videos!**" (05-29). A delivered feature died with NO code change — server-side/device-registration state drift. PC-1..PC-9 are all pipeline-internal; nothing covers environment/external-service drift breaking a shipped feature between two touches. A verification

regime that only re-runs at build/release seams is structurally blind to it (argues for the §11.4.128/§11.4.135 standing-guard direction: guards on a SCHEDULE, not only at seams).

- **NC-2 — Complaint intake latency/evaporation.** The audio-switch-hang ISSUE needed a 13.6 h REMINDER :: re-raise on a different track before an item was minted ~1.6 days after first report. The operator carries the retransmission burden; a session death between report and mint loses the report entirely (the §11.4.210 hook was un-wired precisely then — its own forensic anchor). Intake is a fourth loss seam alongside PC-1/2/6.
- **NC-3 — No operator-checkable liveness signal.** "I toggled the outlet, check 66ff9c4f, however not sure that app is working att all!" (07-16). The operator cannot OBSERVE whether a deployed component is alive; every check routes through the agent, which is exactly the party whose claims are under audit. Deliverables lack an independent, human-readable health surface.

Convergence verdict: the operator's experienced causes CONFIRM PC-2, PC-3, PC-5, PC-6, PC-7, PC-8 independently (different instrument, same physics), add NC-1..NC-3, and provide no evidence against any PC. PC-4 (rules-as-prose) and PC-9 (vocabulary ambiguity) are not directly nameable by a user experiencing failures — their fingerprints are the B/C-class complaints instead.

. Cross-project comparison

Class-A defect reports by project: atmosphere family 18 (track1 12, track3 2, Android-15 old-era 3, main-atmosphere 0 — its window predates the QA drives), claude-toolkit 2, tmux 1, code-server 1, sword-toolkit 1, **helix-code 0, helix-ota 0**.

- Complaints concentrate where a release is DELIVERED TO A TARGET the operator physically touches (device flash, toolkit alias rollout). They are absent where nothing was hand-touched in-window.
- **helix-code (corpus 3, ~0.3% recorded reopens):** 92 operator prompts in-window (denominator needed — 12 of them mention test/release/deploy), zero complaints through the same battery. This is consistent with BOTH "releases work" AND "no first-touch exercise occurred". The toolkit datapoint decides the interpretation: the one time a non-atmosphere deliverable in this corpus WAS first-touched (aliases, 07-20/22), it failed immediately and the operator asked

"wasn't all this tested before the release!?". **Corpus 3's low reopen rate therefore cannot be read as quality evidence; it is detection-pressure-limited** — precisely the Phase-1 hypothesis, now with a live counter-example.

- The complaint SHAPE is uniform across projects (first touch → immediate failure → verification-trust question); only the FREQUENCY differs, and it tracks operator-touch frequency, not project quality.

. Control needles and instrument notes

Instrument incidents hit DURING this mining (all §11.4.201(7) shapes, logged as they happened):

1. **Self-match carrier, twice in one poll.** A liveness poll `pgrep -f mine_complaints.py ... pgrep -f extract_op_prompts` reported both passes RUNNING when **neither process existed** — the poll's own shell cmdline carried both pattern strings. Resolution: `ps -eo pid,args` with the real script path, per §11.4.196(D). Every process-liveness claim below was re-measured with the path-anchored form.
2. **Detached-orphan reaping.** Pass 1 (launched `nohup ... &` inside a background task that exits immediately) was killed externally at ~2.3 GB read with no traceback on stderr — the harness reaps detached children of completed tasks. UNKNOWN: exact reaping trigger. Remedy applied: long passes run as the tracked background command itself (the §11.4.89 pattern); the rerun completed normally.
3. **Truncated-line JSON.** `head -c` sampling produces an unterminated-string parse error on the boundary line by construction — sampling switched to line-whole `head -n / jq`. Malformed-line counts in the real pass are reported in §3, never absorbed.
4. **Over-eager redaction as carrier.** The `sk-...` credential regex matched inside the literal `task-notification ( ...ta + sk-notification... )`, corrupting excerpts. Fixed with a left-boundary guard `(?![A-Za-z0-9])sk-`; §11.4.10 posture retained (over-redaction is the fail-safe direction; no credential is ever reproduced).

5. **Junk classes in the "user" role.** Three harness-generated event classes masquerade as user messages and dominated the naive operator channel: `<task-notification>` blocks (640 of 866 operator-channel matches in the partial pass), compaction summaries ("This session is being continued...", 194/866), and `Caveat:` wrappers. All are excluded by prefix from both the numerator (complaints) and the denominator (operator prompts). A naive "user-role" mining of this corpus without these filters overcounts by >25×.

Control needles:

- N1 (corpus reach): the distinctive operator typo `documetn` (from a real 2026-07-17 interactive directive) found by independent grep in `-run-media... Android-15/93f6aeeb....jsonl` and `-mnt-track1-atmosphere-t1` transcripts — the pipeline's project set demonstrably includes both eras.
- N2 (channel visibility): verified below in §9a after the extract pass — the same event found by grep must appear in the operator-prompt extract.
- N3 (pattern-battery visibility): the known 2026-07-16/17 manual-QA complaint sessions must surface in the complaint set; verified in §4.

§9a — needle results (all run through the SAME pipeline path as the claims they certify):

- N2 PASS: the operator-typed token `documetn` (found by independent grep in raw transcripts) appears in the operator-prompt extract exactly once (`grep -c documetn op_prompts.tsv` → 1) — the extraction path sees real operator events end-to-end.
- N3 PASS: the known 07-16/17 manual-QA complaint sessions surface as 10 + 2 class-A rows with verbatim excerpts (§5, §6) — the battery + ISSUE-prefix union sees the ground-truth channel.
- helix-code absence needed: 92 operator prompts from helix-code passed through the identical `hits_for()` path that produced complaints in 7 other projects; the zero is a measurement, not blindness (§8 carries the honest interpretation limit).
- Battery-miss check (the needle that FAILED and was fixed): 8 of 21
`ISSUE:` -prefixed ground-truth reports carried NO battery pattern ("sign-in as Google Pixel", "brightness sliders do not affect..."), proving phrase-batteries alone under-count first-touch reports by ~35% — the deliverable unions the reporting-directive channel; any future re-run MUST do the same.

- Carrier subtraction verified: 13 prompts carry the pasted §11.4 covenant literal ("...does not work and can[']t be used"), 5 carry the multi-track kickoff mandate ("No change applied using su..."); all 15 are class E, none count as complaints.
- 1 0 Without this subtraction the complaint count inflates ~18% with pure governance boilerplate.

• UNKNOWNs

- **UNKNOWN: everything before 2026-05-26.** Neither store retains earlier transcripts; the April tag-storm era (14 tags in 12 days) has zero operator-voice coverage. The experienced-reopen floor measured here is therefore itself an undercount of the project's full history.
- **UNKNOWN: agent-channel complaint relays.** Pass 1 (which also recorded agent-channel matches) was killed at ~2.3 GB (§9 incident 2); its partial data (1,981 agent rows) was not used. How often the conductor relays operator complaints verbatim into worker prompts is unmeasured.
- **UNKNOWN: the exact reaping trigger** that killed the detached pass-1 process (no traceback, no OOM evidence collected). Operationally closed by running long passes as tracked commands.
- **UNKNOWN: session-file duplication factor.** The same user event appears in multiple session files (forks/resumes — one probe event appeared 6×); hash+day dedup absorbs it for counts, but the per-file multiplicity was not quantified corpus-wide.
- **UNKNOWN: 1:1 mapping of every 0.0.7 QA utterance to a tracker id.** ATM-786/796/798 have no matching `ISSUE:` transcript row (relayed from prose or from a session fragment not captured); no loss PROVEN, but the complete report → item ledger for 07-16 was not established row-by-row.
- **UNKNOWN: whether helix-code deliverables work** (§8) — zero complaints is consistent with both readings; only a first-touch exercise decides it.
- **PENDING_FORENSICS: NC-1 scope.** How many other "worked yesterday, same image" failures exist is unmeasured — the battery has no pattern for external-state drift phrased positively; only the 05-29 instance was found (via `same_again`).

. Deliverables and reproduction

- `complaints.tsv` (beside this file): 91 classified, redacted, deduped operator-channel rows — columns `ts`, `class`, `store`, `project`, `session`, `branch`, `patterns`, `strong`, `subsystems`, `hash`, `textlen`, `excerpt` .
- Class key: `A_defect_report` (verified first-touch defect),
`B_recurrence_meta` , `C_verification_trust` , `D_qa_pressure` /
`D_directive_context` (not complaints; context), `E_carrier_*` (false positives, retained for audit).
- Reproduction: `/tmp/opm/{extract_op_prompts.py, classify_op.py, final_classify.py}` (session-local; re-runnable read-only against the stores). Every count in this report is regenerated by those three scripts.
- §11.4.10: all excerpts passed THREE redaction layers (credential-shape, e-mail/ `пароль` , and a post-write high-entropy-token scrub after one `su` credential was found surviving the phrase-shaped regexes — found by audit, scrubbed, and the residual scan now shows only the benign "ri sk- free"
1 2 carrier). The QA-account prompt (Qobuz) is redacted in the deliverable.

. Anti-bluff certification

No count in this document is estimated: each traces to a command over the extract or the read-only tracker DB. Absences carry needles (§9a). The two instrument failures that occurred during the work (self-match carrier, detached-orphan reaping) are reported in §9, not absorbed. The corpus truncation boundary (2026-05-26) is stated wherever it bounds a claim. Transcripts were opened READ-ONLY; nothing under `~/claude*` was modified; the tracker DB was opened `mode=ro` ; no commit/push was performed.