

Helix Constitution “ Universal GEMINI.md

Field	Value
Revision	3
Created	2026-06-09
Last modified	2026-06-21T00:00:00Z
Status	active
Status summary	Added Â§11.4.163 mirror “ Universal Media Validation: every recorded artifact MUST pass MEDIA VALIDATION pipeline (OCR/transcription/text parsing vs patterns, self-validated analyzer, structured verdict, pinpoint on FAIL, real-time trigger, Â§1.1 mutation). Added Â§11.4.164 mirror “ Universal Auto-Propagation Hook: post_update_hook.sh detects/registers/installs skills/MCP/hooks/scripts; consumer invokes after every pull. Added Â§11.4.165 mirror “ Universal Independent Verification Agent: every output passes INDEPENDENT verifier (Â§11.4.70/.20), zero-finding GO. All Classification: universal (Â§11.4.17). Prior round: Added Â§11.4.160“162 mirror entries “ Vision-verified recording + HelixQA bridge (User mandate 2026-06-21): every video recording MUST be processed through a vision/OCR pipeline reading on-screen content against SPECIFY-phase expected patterns; 5s frame capture, self-validated analyzer, per-frame PASS/FAIL with evidence path. Rootless container runtime (User mandate 2026-06-21): Podman rootless mandatory; rootful Docker/sudo FORBIDDEN unless Â§11.4.112; vasic-digital/containers sole orchestration layer. OpenDesign UI design system (User mandate 2026-06-21): mandatory UI design-and-refinement system for all user-facing interfaces; design tokens for color/typography/spacing; light+dark themes; extend upstream per Â§11.4.74. All three Classification: universal (Â§11.4.17). GEMINI.md mirrors updated in lockstep per Â§11.4.157. Prior round: New Gemini CLI carrier of the Helix Constitution “ added by Â§11.4.140 (universal action-prefix system) so Gemini CLI gets the LAYER-1 action-prefix recognition floor out of the box, plus the Â§11.4.141 token-efficiency mandate mirror. This file is one of the four canonical agent context carriers (CLAUDE .md,

Field	Value
	AGENTS.md, QWEN.md, GEMINI.md) the constitution maintains per Â§11.4.35; the textual identity of the action-prefix recognition block across all four IS the LAYER-1 universality. This GEMINI.md is a thin carrier â€” the full universal rule set is in Constitution.md (canonical) and the sibling carriers; Gemini CLI reads this file as its context carrier. Back-filled Â§11.4.142â€”157 to lockstep (Â§11.4.157).
Issues	none
Issues summary	â€”
Fixed	Â§11.4.140 + Â§11.4.141 mirrors
Fixed summary	Carrier created in lockstep with the Constitution.md Â§11.4.140 + Â§11.4.141 additions.
Continuation	â€”

How inheritance works

This GEMINI.md is the **Gemini CLI carrier** of the Helix Constitution submodule. It is one of the four canonical agent context carriers (constitution/CLAUDE.md, constitution/AGENTS.md, constitution/QWEN.md, constitution/GEMINI.md) â€” these ARE the canonical root per Â§11.4.35, NOT consumer extensions. The full universal rule set lives in constitution/Constitution.md; the sibling carriers (CLAUDE.md / AGENTS.md / QWEN.md) carry the same content for their respective agents. A consuming projectâ€™s own repo-root GEMINI.md (if any) opens with the inheritance pointer and adds only project-specific rules; THIS file is the canonical carrier and does not.

For agents driven from a consuming project, the projectâ€™s repo-root carrier SHOULD start with a clearly-marked inheritance pointer:

```
## INHERITED FROM constitution/GEMINI.md
```

All rules in `constitution/GEMINI.md` (and the `constitution/Constitution.` it references) apply unconditionally. The project-specific rules below extend them.

Action-prefix recognition (Â§11.4.140 â€” LAYER-1)

The block below is the **canonical LAYER-1 action-prefix recognition instruction**, embedded VERBATIM (identical text) into every agent context carrier the constitution maintains. That textual identity across all four carriers IS the LAYER-1 universality: Gemini CLI reads this block via this carrier and applies the prefix expansion itself, so the action-prefix system works on Gemini CLI with zero extra setup.

Â§11.4.140 â€” Universal action-prefix system (ACTION_NAME : :) (User mandate, 2026-06-09; GRAMMAR_ADDENDUM 2026-06-09). When a user promptâ€™s FIRST non-blank line starts with a recognised action prefix, you MUST: (1) look the action token up in the action registry constitution/

actions/registry.yaml (or \$HELIX_ACTION_REGISTRY); (2) if it is a registered action, REPLACE the prefix with that action's expansion text and apply its rules; (3) execute the remainder of the prompt under the expanded instruction. **Four EQUIVALENT forms** " same action, same expansion, same execution: (1) ACTION_NAME :: <rest> (bare ::), (2) PREFIX::ACTION_NAME :: <rest> (namespaced ::), (3) /ACTION_NAME <rest> (bare slash), (4) /PREFIX::ACTION_NAME <rest> (namespaced slash). Thus BACKGROUND :: x %j; DEFAULT::BACKGROUND :: x %j; /BACKGROUND x %j; /DEFAULT::BACKGROUND x. PREFIX is an action NAMESPACE; the reserved default namespace is DEFAULT, and an action runs WITH or WITHOUT the prefix. Grammar (all hold): anchored at the FIRST non-blank line only (mid-prose tokens never match); the action token AND the namespace are UPPERCASE-only [A-Z] [A-Z0-9_]* (lowercase never matches); the namespace separator :: inside the token carries NO surrounding spaces (PREFIX::ACTION_NAME), DISTINCT from the action-body separator " :: " (one ASCII space on each side of :: " avoids C++ Foo::Bar, YAML key: value, URLs) in forms 1/2 and the slash-body separator (one space) in forms 3/4; stacked prefixes (A :: B :: rest) apply outer-to-inner, left-to-right (expand A, re-scan, expand B, then the residual is the task); a leading \ escapes the prefix for BOTH the :: and the slash form (\BACKGROUND :: x, \ / BACKGROUND x " treat literally, strip the backslash, NO expansion) so action names can be discussed. **Conflict rule (slash form):** /ACTION_NAME (form 3) is honored as the action ONLY when ACTION_NAME (case-folded) does not collide with a built-in/host slash command (registry slash_bare: auto + slash_conflicts: [...]); form 4 (/PREFIX::ACTION_NAME) is ALWAYS unambiguous and always honored. An unknown token that matches the grammar shape (any of the 4 forms) but is NOT registered is NEVER silently expanded or silently dropped " ask which registered action was meant (§11.4.66 / §11.4.105) or treat it as a literal prompt, NEVER invent an expansion (§11.4.6); any prompt not satisfying the grammar is an ordinary prompt and the system is a no-op. The registered action BACKGROUND expands to: *“The following prompt that we will provide MUST BE executed in background in parallel with all main work streams using the subagents-driven development approach! All work done MUST PRODUCE rock solid evidence covered with hard physical proof(s) that all done is working as expected and as specified without any false results and without any bluff!”* (composes §11.4.20 / §11.4.70 subagent-driven, §11.4.58 / §11.4.103 parallel streams, §11.4.89 background execution, §11.4.5 / §11.4.69 / §11.4.107 captured physical evidence, §11.4 anti-bluff). The system is UNIVERSAL (every CLI agent reads this block via its context carrier per §11.4.35), extensible (new action = new registry row), decoupled + reusable (§11.4.28), and loads out-of-the-box. Classification: universal (§11.4.17). **Canonical authority:** constitution submodule [Constitution.md](#) §11.4.140. Non-compliance is a release blocker. No escape hatch " no --skip-action-prefix, --ignore-prefix, --no-registry, --invent-expansion-OK, --single-layer-only flag.

Token-efficiency mandate (§11.4.141)

§11.4.141 " Token-efficiency mandate (research-derived + operator mandate, 2026-06-09). Every project worked on by AI coding agents MUST cut token spend (input AND output) toward 30% of current (a 60%–70% reduction) WITHOUT degrading quality/performance/safety or breaking any existing mechanism, via a composable, safety-ranked measure set: (1) **prompt-cache the static governance prefix** " the always-loaded governance

forms a byte-stable cache-breakpointed prefix with no volatile bytes ahead of it; cache reads cost ~0.1Å— base input (the dominant cost driver â€” measured ~170K tokens of governance re-sent every turn, externally corroborated by Claude Code issue #24147); caching is transparent so it removes no rule, weakens no gate, changes no verdict â€” only billing (PRIMARY, biggest + safest lever); the operative discipline is to **batch governance edits and keep CLAUDE.md + the constitution byte-stable mid-session** â€” every governance edit busts the prompt cache (BASELINE.md finding); (2) **subagent model-tiering + output-to-file** â€” mechanical non-judgment work (search/grep/status/doc-export/read-only probes) to a Haiku-class model, the strong model RESERVED for all reasoning/verdicts/fix-design (Â§11.4.102)/code-review (Â§11.4.125)/demotion (Â§11.4.7), large output persisted to a file not an inline 350â€”520K-token transcript; the cheap model never emits a PASS so Â§11.4.50 + anti-bluff are untouched; (3) **thin always-loaded INDEX + on-demand detail** â€” concise index (one line per fix/anchor, EACH carrying the literal 11.4.N token so propagation gates pass) with the canonical full text kept gate-scanned in constitution/Constitution.md and reachable in one hop â€” a de-duplication realising Â§11.4.35, never a deletion; (4) **CodeGraph/retrieval-first over full-file loading** (Â§11.4.78/Â§11.4.79/Â§11.4.80); (5) **output-token reduction** â€” terse status + effort: "low" on the mechanical allowlist only; (6) **tool-call batching + no re-reads**; (7) **compaction/context-editing for long sessions. Mandatory measured proof:** a token-accounting harness measures tokens-per-development-cycle BEFORE vs AFTER on a frozen deterministic workload from the authoritative usage object (input/cache_read/cache_creation/output split; NEVER tiktoken, NEVER the client-side cost estimate), reproduced N times (Â§11.4.50), pass = AFTER â‰¥ 40% of BEFORE OR the measured best-safe reduction with a cited cold-cache reason; the AFTER run MUST show ZERO regression on the pre-build sweep + meta-test mutation sweep + propagation gates + a strong-model reasoning probe + a cache-warm proof (cache_read_input_tokens > 0) â€” cost reduction with quality regression is a Â§11.4 FAIL. The headline number is the *measured* reduction, never the design estimate (Â§11.4.6/Â§11.4.123). No measure may break/degrade any existing mechanism, and the rule is structured so none can. Composes

Â§11.4.5/.6/.20/.40/.50/.58/.69/.70/.78/.79/.80/.103/.106/.123/.125/.128/Â§12.6/Â§1.1.

Classification: universal (Â§11.4.17). Propagation gate CM-COVENANT-114-141-PROPAGATION (literal 11.4.141) + recommended gate CM-TOKEN-EFFICIENCY + paired Â§1.1 mutation (inject a pre-breakpoint volatile token â†’ cache collapses â†’ measured reduction falls below bar â†’ gate FAILs). **Canonical authority:** constitution submodule [Constitution.md](#) Â§11.4.141. Non-compliance is a release blocker. No escape hatch â€” no --skip-token-efficiency, --no-cache-governance, --assert-reduction-without-measuring, --tier-down-reasoning, --inline-all-governance, --tiktoken-estimate-OK flag.

Â§11.4.142 â€” Universal code-review mandate â€” every change reviewed, always, no exception (User mandate, 2026-06-09). Verbatim operator mandate: â€œALL changes we do MUST pass through the code review step!!! ALWAYS!!!â€” EVERY change made to ANY repository this Constitution governs â€” without exception â€” MUST pass through an INDEPENDENT code-review step BEFORE it is accepted, committed, or built. NO change class is exempt: source code, fixes, tests, gates, meta-test mutations, documentation, doc-tooling, build/CI scripts, configuration, governance files (Constitution / CLAUDE / AGENTS / QWEN), conductor main-stream edits, sub-agent output, refactors, single-line edits â€” if a diff exists, it gets an independent review. This is the ABSOLUTE form of Â§11.4.125 (code-review agent gate after a batch, before pre-build sweep + main build): Â§11.4.142 strips every implicit scoping Â§11.4.125â€™s â€œafter all fixes/changes/implementations are doneâ€” phrasing could leave open â€” no â€œjust a doc editâ€”, no â€œjust a one-linerâ€”, no â€œthe author already self-reviewed (Â§11.4.92)â€”, no â€œtrivial changeâ€” carve-out. **Independence is load-bearing** â€” the reviewer MUST be structurally separated from the author (a dedicated code-review agent, subagent-driven by default per Â§11.4.70 / Â§11.4.20, or a distinct human), NEVER the author re-reading their own work; Â§11.4.92â€™s multi-pass self-evaluation is the AUTHOR-side

discipline and PRECEDES (never satisfies) Â§11.4.142. **The review is itself anti-bluff** (Â§11.4 / Â§11.4.1) â€” a rubber-stamp â€œlooks goodâ€” is not a review; it MUST genuinely analyse correctness, safety (no host Â§12 / data Â§9 / target-hardware Â§11.4.133 regression), will-it-really-work (no solve-A-create-B), end-user behaviour (Â§11.4 / Â§107), and test-genuineness (Â§1.1), its conclusions captured evidence per Â§11.4.5 / Â§11.4.69, and **it iterates to a clean GO per Â§11.4.134** â€” any finding (BLOCKING / nit / warning) re-arms the loop, acceptance only on ZERO new findings + ZERO warnings with rock-solid physical evidence. Honest boundary (Â§11.4.6): a passing review does NOT replace Â§11.4.108 four-layer runtime-signature verification nor the Â§11.4.40 full-suite retest â€” it is one of MULTIPLE STRONG LAYERS, and the FIRST one every change crosses. Composes Â§11.4.1 / Â§11.4.4 / Â§11.4.5 / Â§11.4.6 / Â§11.4.20 / Â§11.4.40 / Â§11.4.69 / Â§11.4.70 / Â§11.4.92 / Â§11.4.108 / Â§11.4.110 / Â§11.4.125 / Â§11.4.134 / Â§107 / Â§1.1. Classification: universal (Â§11.4.17) â€” the consuming project supplies its reviewer-dispatch mechanism + change-acceptance seam (commit wrapper / merge queue / PR gate) per Â§11.4.35. Propagation gate CM-COVENANT-114-142-PROPAGATION (literal 11.4.142) + recommended gate CM-EVERY-CHANGE-REVIEWED (every accepted change carries a fresh independent-review-completed marker for its diff, produced by a reviewer distinct from the author, before acceptance/commit/build) + paired Â§1.1 mutation (strip the literal â†’ propagation gate FAILs; accept a diff with no independent-review marker â†’ CM-EVERY-CHANGE-REVIEWED FAILs). **Canonical authority:** constitution submodule [Constitution.md](#) Â§11.4.142. Non-compliance is a release blocker. No escape hatch â€” no --skip-review, --no-review, --trivial-change-exempt, --doc-edit-exempt, --self-review-suffices, --review-after-commit flag.

Â§11.4.143 â€” Real-user-journey mandate for video-streaming-app full-automation tests (User mandate, 2026-06-10). Verbatim operator mandate: â€œAll video streaming apps â€” require to choose some title and to press proper UI button to start or resume playing! Proper UI interaction to play exact show with proper content and subtitles is MANDATORY! Without it we just eventually play on 2nd display sample, and thatâ€™s it mostly! THIS MUST BE ADDED as MANDATORY RULE regarding testing of any video streaming app in general with full automation tests! Universal in root constitution + ATMOSphere project extensions.â€” Any full-automation test that asserts a video player / streaming application plays content MUST drive the REAL end-user journey through the appâ€™s OWN UI â€” launch â†’ BROWSE the actual catalog â†’ choose a SPECIFIC title â†’ press the real Play/Resume button â†’ confirm THAT chosen content is genuinely playing, with its correct subtitles, on the intended routing target. A test that bypasses the journey with a sample / demo / built-in-loop clip, a deep-link / am_start - a VIEW / intent shortcut, a synthetic or pre-staged stream, or any path that does NOT exercise the appâ€™s own browse-select-play UI is a Â§11.4 PASS-bluff at the user-journey layer: it validates ROUTING (that *something* reaches the display) while leaving the user-visible behaviour the operator cares about â€” â€œthe show I picked actually playsâ€” unproven (the operatorâ€™s â€œwe just eventually play on 2nd display sample, and thatâ€™s it mostlyâ€” gap). The mandate (ALL hold): (1) **Real journey, not a shortcut** â€” launch â†’ catalog browse â†’ specific-title selection â†’ real Play/Resume press through the appâ€™s own UI (Â§11.4.48 UI-driven), NEVER a deep-link / intent / sample / loop-clip shortcut; (2) **Chosen-content confirmation** â€” the PASS proves the SPECIFIC selected title is playing (not merely that pixels move on the target) via the Â§11.4.107 liveness battery on the Â§11.4.136 real-content path + the Â§11.4.137 subtitle content-correctness oracle for the chosen titleâ€™s captions; (3) **Non-introspectable UIs use the pixel oracle** â€” when the appâ€™s accessibility hierarchy is blank/unreliable (TV-Compose / leanback / canvas / GL), DRIVE input + ASSERT content via the Â§11.4.117 CV/OCR pixel oracle, never a hierarchy-only tool; (4) **Login via the credential single-source** (Â§11.4.10), never hardcoded, never logged; (5) **Honest SKIP, never a faked PASS** â€” where the autonomous journey is genuinely infeasible (hard human-only login / CAPTCHA, geo-block per Â§11.4.3, secure-surface pixel-blanking per Â§11.4.112) the test is operator_attended SKIP-with-reason per Â§11.4.52 + Â§11.4.3 with a tracked migration

item â€” NEVER a metadata-only / sample-played / routing-only PASS. Honest boundary (Â§11.4.6): â€œthe routing fired so the title is playingâ€” is a guess â€” only the chosen-content liveness + subtitle oracle on the real journey proves it. Composes Â§11.4.48 / Â§11.4.107 / Â§11.4.117 / Â§11.4.136 / Â§11.4.137 / Â§11.4.52 / Â§11.4.3 / Â§107 / Â§1.1. Classification: universal (Â§11.4.17) â€” the consuming project supplies its concrete app roster, login-credential source, routing target, and UI-driving / pixel-oracle harness per Â§11.4.35. Propagation gate CM-COVENANT-114-143-PROPAGATION (literal 11.4.143) + recommended gate CM-VIDEO-REAL-JOURNEY-TEST (every video-streaming-app playback test drives the real browse-select-play journey + confirms the chosen content + subtitles, or SKIPs-with-reason) + paired Â§1.1 mutation (replace a real-journey testâ€™s browse-select-play path with a deep-link / sample shortcut â†’ gate FAILs; strip the literal â†’ propagation gate FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule [Constitution.md](#) Â§11.4.143. Non-compliance is a release blocker. No escape hatch â€” no --sample-playback-ok, --skip-real-journey, --deep-link-suffices, --routing-only-pass, --no-title-selection flag.

Â§11.4.144 â€” Tracked/recorded-device availability-following mandate (User mandate, 2026-06-10). Direct operator mandate: every device the project is tracking / following / recording (test / debug / manual-testing device, across every reachable transport â€” USB / wireless ADB / SSH / serial / network introspection API) MUST be availability-FOLLOWED â€” its connection state continuously monitored, any drop handled, never silently abandoned and never presented as a continuous recording. The Â§11.4.128 always-on recorder KNOWS a tracked device is absent (its per-device loop guards on the reachable state) but, lacking a following discipline, merely spins idle â€” no data captured, no offline event logged, no resume, no escalation â€” so the recording corpus presents a continuous timeline with a silent hole, a Â§11.4 PASS-bluff at the recording-integrity layer (it claims continuous capture while no data exists for the gap). On a tracked device leaving its reachable state the system MUST, automatically: (a) **DETECT** the drop and **log an honest offline event** into the recording corpus (Â§11.4.6 â€” a silent gap presented as continuous capture is a fabricated-continuity bluff; Â§11.4.128 â€” a silent recording gap defeats always-on recording); (b) **WAIT** for the device to return using the projectâ€™s ALREADY-DEFINED reconnection timings â€” never invented numbers (Â§11.4.6), the SAME grace / reconnect / poll budgets the projectâ€™s recovery path already uses; (c) **RE-ATTACH** â€” resume recording / tracking the moment the device returns and log an honest online / resume event; (d) **ESCALATE** to the projectâ€™s sanctioned device-recovery path (per Â§11.4.69 feature class `device_recovery`) if the device does not return within the defined timeout â€” through the sanctioned, authorization-gated recovery entry point ONLY, never bypassing its gate and never performing a destructive recovery (e.g. a power-cycle) autonomously without that authorization (Â§11.4.21 / Â§11.4.101 â€” high-blast-radius recovery is gated; while blocked the system keeps following the device, logging the blocked-escalation honestly). A tracked-device drop that produces a silent corpus hole, a never-resumed recording, or a never-escalated permanent absence is the bluff this anchor forbids. Composes Â§11.4.128 (always-on recording â€” Â§11.4.144 closes its drop-handling gap) / Â§11.4.69 (`device_recovery` sink-side positive evidence) / Â§11.4.6 (honest offline / online events, reused-not-invented timings) / Â§11.4.14 (watchdog children reaped on stop) / Â§11.4.21 + Â§11.4.101 (gated, non-autonomous escalation). Classification: universal (Â§11.4.17) â€” the consuming project supplies its concrete tracking transport, device set, already-defined reconnection timings, and sanctioned recovery entry point per Â§11.4.35. Propagation gate CM-COVENANT-114-144-PROPAGATION (literal 11.4.144) + recommended gate CM-DEVICE-AVAILABILITY-FOLLOWED + paired Â§1.1 meta-test mutation (strip the watchdog wiring / let an absence go unlogged â†’ gate FAILs; strip the literal â†’ propagation gate FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule [Constitution.md](#) Â§11.4.144. Non-compliance is a release blocker. No escape hatch â€” no --skip-availability-following, --silent-recording-gap-OK, --no-reconnect-wait, --invent-

reconnect-timing, --skip-recovery-escalation, --autonomous-power-cycle-OK flag.

Â§11.4.145 â€” Independent multi-angle impact-research per change (User mandate, 2026-06-10). For EVERY fix / change / new feature, INDEPENDENT impact-research agents (subagent-driven Â§11.4.70/Â§11.4.20, structurally separate from the author â€” Â§11.4.92 self-eval PRECEDES, never satisfies â€” adversarial â€œrefute-the-changeâ€” stance to defeat the documented LLM confirmation-bias failure mode) MUST research the change AND its connected/dependent features across a CLOSED SET OF EIGHT ANGLES â€” (1) correctness/logic, (2) regression (what existing working feature could break â€” via call-graph impact enumerating every direct + transitive caller and contract-dependent feature, each mapped to its test), (3) latent/dangerous-code â€” the recents class (runtime-only failure, interface/ABI/contract mismatch, concurrency/data-race, resource lifecycle), (4) security (taint/injection/secret/unsafe-API), (5) performance (latency/memory/thermal under load vs baseline), (6) host/data/target-hardware safety per Â§12/Â§9/Â§11.4.133, (7) cross-feature interaction (shared state/timing/hardware contention), (8) business-logic conformance (matches the spec AND the connected featuresâ€™ contracts, not merely compiles). Forensic anchor (FACT, project-internal): the recents / 3-button-nav path shipped a latent AIDL-interface-contract mismatch that PASSED code review yet failed at RUNTIME ONLY (ANR on a recents-reaping gesture) because no review angle interrogated the interface contract / concurrency-lifecycle / silently-broken connected feature â€” the Â§11.4.108 SOURCEâ€™ARTIFACTâ€™RUNTIMEâ€™USER-VISIBLE gap caught only after the fact; Â§11.4.145 shifts that discovery LEFT. Each angle MUST name the TOOL(S) used + obtain the required DEPTH (the diff, the full changed unit, its declared contract, the spec section, every connected feature â€” NEVER the diff lines alone) + emit a captured-evidence conclusion per Â§11.4.5/Â§11.4.69 (â€œlooks fineâ€” forbidden, Â§11.4.1); a genuinely-N/A angle is recorded NOT_APPLICABLE: <reason> per Â§11.4.6, never silently skipped. Output = a per-change impact-research REPORT (one section per angle: tool + evidence path + verdict) + a single GO/NO-GO that BLOCKS acceptance/commit/build on ANY unmitigated risk in ANY angle; the change is fixed/mitigated + affected angles re-researched, iterating to a CLEAN GO (zero unmitigated risk + zero warning) per Â§11.4.134 before the Â§11.4.125 review gate. Honest boundary (Â§11.4.6): a GO proves cross-angle internal-consistency + bounded blast-radius â€” it does NOT replace Â§11.4.108 runtime-signature verification nor Â§11.4.40 full-suite retest; it is one of MULTIPLE STRONG LAYERS, the FIRST research pass every change crosses. Composes/strengthens Â§11.4.92/.125/.108/.110/.134/.78/.79/.107/.85/.133/Â§12/Â§9/.99/.6/Â§1.1. Classification: universal (Â§11.4.17). Propagation gate CM-COVENANT-114-145-PROPAGATION (literal 11.4.145) + recommended gate CM-IMPACT-RESEARCH-PER-CHANGE + paired Â§1.1 meta-test mutation (strip the literal â€™ propagation gate FAILs; accept a change with a report missing an angle or an unmitigated NO-GO â€™ CM-IMPACT-RESEARCH-PER-CHANGE FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule [Constitution.md](#) Â§11.4.145. Non-compliance is a release blocker. No escape hatch â€” no --skip-impact-research, --single-angle-suffices, --author-researches-own-change, --diff-skim-OK, --no-go-overridable, --trivial-change-no-research flag.

Â§11.4.146 â€” Reproduce-first test + same-test-confirms-fix + mandatory extend-to-all-cases workflow (User mandate, 2026-06-10). Every reported problem MUST be handled by a NAMED three-step test workflow â€” Â§11.4.146 does NOT re-author its component disciplines, it NAMES + BINDS them into the operatorâ€™s exact per-defect sequence and adds ONLY the two emphases the components leave implicit. **(STEP 1 â€” REPRODUCE-FIRST + INVESTIGATE)** before any fix, author the Â§11.4.43 RED test as a Â§11.4.115 RED-baseline-on-the-broken-artifact (reproduce the defect on the CURRENT pre-fix artifact, capture defect-present physical evidence per Â§11.4.5/Â§11.4.69/Â§11.4.107); NEW EMPHASIS (D1) that same RED test is ALSO a deliberate investigation instrument per Â§11.4.102(A) â€” it MUST

gather ADDITIONAL forensic data characterising the defect (triggers, boundaries, input/topology scope, adjacent failure modes) that feeds BOTH the fix design AND the STEP-3 extend scope; a RED test used only as a binary present/absent gate with no characterisation satisfies Â§11.4.115 but NOT Â§11.4.146 STEP 1. **(STEP 2 â€” SAME-TEST-CONFIRMS-FIX)** after the fix the SAME test source (the Â§11.4.115 polarity switch flipped RED_MODE=1 â† ’ 0) confirms the defect ABSENT â€” RED-on-broken then GREEN-on-fixed, both captured, on a clean target per Â§11.4.108/Â§11.4.139, validated-first-after-redeploy per Â§11.4.130, deterministic per Â§11.4.50; no separate happy-path test substitutes (the Â§11.4.43/Â§11.4.115 PASS-bluff). **(STEP 3 â€” EXTEND-TO-ALL-CASES, mandatory per-fix)** NEW EMPHASIS (D2) immediately after STEP 2 the test MUST be FANNED OUT across the full case-space of the SAME functionality â€” all flows, valid + invalid + boundary (Â§11.4.85 empty/max/off-by-one) + concurrent (Â§11.4.85 contention) + failure-injection (Â§11.4.85 chaos) + topology variants (Â§11.4.3) â€” confirming no issue of any kind, with PROVABLE enumerated coverage per Â§11.4.118 (a listed case-set with per-case outcome, never â€œno other issues foundâ€”); a REQUIRED step, NOT deferred to a release-cycle discovery sweep and NOT reduced to a single guard; each user-visible case carries rock-solid physical evidence per Â§11.4.123 + is registered into the Â§11.4.135 standing regression-guard suite; a newly-discovered fan-out defect triggers Â§11.4.4 test-interrupt + re-enters STEP 1. (ANTI-BLUFF, all steps) every PASS is rock-solid CAPTURED physical evidence per Â§11.4.123 (Â§11.4.5/Â§11.4.69/Â§11.4.107) â€” metadata-only / config-only / absence-of-error / grep-without-runtime PASS forbidden (Â§11.4/Â§11.4.1); unclear validation method â† ’ deep-research-before-declaring-untestable per Â§11.4.123/Â§11.4.8/Â§11.4.99; an operator-found defect the green suite missed triggers Â§11.4.138. Honest boundary (Â§11.4.6): STEP 3 reduces the unknown-unknown surface but does not prove zero remaining defects (Â§11.4.118) â€” un-exercised cases stated as honest gaps, never silently implied clean. Composes Â§11.4.43/.115/.102/.130/.108/.139/.50/.85/.118/.135/.3/.123/.5/.69/.107/.138/.4/Â§107/Â§11.1. Classification: universal (Â§11.4.17). Propagation gate CM-COVENANT-114-146-PROPAGATION (literal 11.4.146) + recommended gate CM-REPRODUCE-FIRST-THEN-EXTEND (every closed defectâ€™s fix carries a Â§11.4.115 reproduce-first polarity test with captured defect-characterisation [STEP 1] + its RED_MODE=0 GREEN confirmation [STEP 2] + an enumerated per-functionality extend case-set registered into the Â§11.4.135 suite [STEP 3]; a closure with the reproduceâ€™confirm pair but NO enumerated extend case-set FAILs) + paired Â§1.1 meta-test mutation (strip the literal â† ’ propagation gate FAILs; close a defect with only the reproduceâ€™confirm pair and no extend-case-set â† ’ CM-REPRODUCE-FIRST-THEN-EXTEND FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule [Constitution.md](#) Â§11.4.146. Non-compliance is a release blocker. No escape hatch â€” no --skip-reproduce-first, --fix-without-red, --skip-extend-to-all-cases, --reproduce-confirm-suffices, --defer-extend-to-release-sweep, --single-guard-suffices, --no-defect-characterisation flag.

Â§11.4.147 â€” Crashed-agent respawn-until-complete + no-work-loss registry mandate (User mandate, 2026-06-10). Verbatim operator intent: â€œany agent that crashed because of something MUST BE respawned and finish its work at some point! We MUST NOT lose any work, forget about or have it corrupted!â€” Forensic case study (FACT): dispatched subagents died on TRANSIENT causes (API Error: Server is temporarily limiting requests rate-limit killed 5 at once, API Error: socket connection closed unexpectedly killed 2, one left PARTIAL edits â€” two source files written, the dependent SystemUI sliders + doc + test NOT done), each re-dispatched by pure conductor vigilance with NO mechanical guarantee â€” the moment conductor context is lost / compacted / the conductor itself crashes, an in-flight-but-dead work unit is silently dropped (lost / forgotten / corrupted). Every dispatched agent/subagent MUST be tracked through its full lifecycle so a crash NEVER loses, forgets, or corrupts its work; a crashed agent is NOT a completed agent (Â§11.4.6 â€” â€œit probably finished before dyingâ€” is a forbidden guess; abnormal termination is positive evidence the unit is OPEN). The mandate (ALL hold): **(a) durable agent REGISTRY** â€” every

dispatched agent has an append-only machine-readable registry entry (stable id, task + Â§11.4.58 file-scope, declared output path(s) + Â§11.4.5/Â§11.4.69 evidence dir, dispatch timestamp, status from the CLOSED SET {dispatched | in-flight | crashed | respawned | complete}), reusing the Â§11.4.116 sync substrate (JSONL event stream + atomic status snapshot; â€œan entry with no dispatch-event cannot show completeâ€œ); the registry is the SINGLE SOURCE OF TRUTH for â€œis this work owed?â€œ and an unregistered dispatch is itself a violation; (b) **mechanical CRASH-DETECTION + RESPAWN-until-complete** â€œany abnormal termination (transient rate-limit/socket-close, any non-completion exit, or a terminal error after the runtimeâ€™s own retries are exhausted) flips the entry to crashed + keeps the unit OPEN, and the unit is respawned (fresh agent claims the same id + scope + preserved state) and re-responded until an agent reaches complete; respawn is the safe/reversible/bounded decision taken autonomously per Â§11.4.101 (never blocks the loop for a human), transient causes use the projectâ€™s ALREADY-DEFINED backoff budgets (never invented, Â§11.4.6), a deterministically-reproducible non-transient terminal error is investigated per Â§11.4.102 before further respawn (never a blind retry loop); (c) **PARTIAL-STATE preserve + Â§11.4.84-check + resume-or-clean-restart** â€œthe crashed agentâ€™s uncommitted edits + output doc are PRESERVED (never silently discarded + lost, never blindly committed + corruption), the respawn runs the Â§11.4.84 quiescence check on the preserved tree (every modified file accounted-for vs scope, no mutation// always pass/_mutated_* residue, no half-written/torn artifact), then EITHER RESUMES idempotently when it passes OR CLEAN-RESTARTS from a known-good base (Â§9.2 pre-op backup, reversible Â§11.4.101) when inconsistent â€œnothing lost, nothing corrupted; per-agent git worktree isolation (Â§11.4.58 L4 / Â§11.4.84) keeps the partial tree from contaminating other streams; (d) **crash + done + COMPLETION criterion** â€œa unit is DONE only when an agent reaches complete with its required captured evidence/output landed (Â§11.4.5/Â§11.4.69 + the Â§11.4.116 verdict-carries-evidence-path rule); the endless-loop done-condition (Â§11.4.87/Â§11.4.94/Â§11.4.97/Â§11.4.126) MUST NOT read satisfied while any entry is dispatched/in-flight/crashed/respawned-not-yet-complete, the zero-idle survey (Â§11.4.94) treats every non-complete entry as an OPEN item to reclaim, and a registry showing complete without landed evidence is a Â§11.4 PASS-bluff at the agent-lifecycle layer. Honest boundary (Â§11.4.6): the registry + respawn guarantee work is not lost/corrupted, NOT that it is correct â€œthe respawned output still crosses Â§11.4.108 + Â§11.4.125/Â§11.4.142 review + Â§11.4.40 retest; Â§11.4.147 is the durability layer beneath those, the agent-side analogue of Â§11.4.144 (same detect+wait/backoff+re-attach/respawn+escalate shape). Composes Â§11.4.6/.58/.84/.87/.94/.97/.101/.116/.102/.108/.125/.142/.40/.128/.144/Â§9.2/Â§1.1. Classification: universal (Â§11.4.17). Propagation gate CM-COVENANT-114-147-PROPAGATION (literal 11.4.147) + recommended gate CM-CRASHED-AGENT-RESPAWN-TRACKED + paired Â§1.1 meta-test mutation (strip the literal + propagation gate FAILS; mark a crashed entry as the loopâ€™s done-condition complete without landed evidence, OR drop a dispatched agent without a registry entry + CM-CRASHED-AGENT-RESPAWN-TRACKED FAILS; gate-code = separate work item). **Canonical authority:** constitution submodule [Constitution.md](#) Â§11.4.147. Non-compliance is a release blocker. No escape hatch â€œno --skip-agent-registry, --crash-equals-done, --no-respond, --discard-partial-state, --blind-commit-partial, --forget-dead-agent, --loop-done-with-crashed-entries flag.

Â§11.4.148 â€œWorkable-item integrity (status+type+id) + comprehensive structured description + bidirectional external-tracker sync + **BLOCKED unblock-choices mandate (User mandate, 2026-06-10)**. BINDS + STRENGTHENS the workable-item discipline (Â§11.4.15 status / Â§11.4.16 type / Â§11.4.54 id / Â§11.4.21 Operator-blocked / Â§11.4.91 description-clarity / Â§11.4.93 + Â§11.4.95 SQLite SSoT / Â§11.4.106 docs_chain) into ONE integrity contract spanning DB + docs + external tracker, adding the operatorâ€™s three emphases: (D1) **no item without a valid status + valid type + stable id** â€œon ALL three surfaces (an item missing any of the three in the DB, the rendered docs, OR the external tracker

FAILs the validator “release-blocker; the id is the cross-surface binding key); **(D2) comprehensive structured description per item** “WHAT it is (Â§11.4.91 “6-word/40-char clear meaning) + HOW it manifests + HOW to reproduce (Â§11.4.115/.146) + ACCEPTANCE CRITERIA (Â§11.4.123/.69 captured-evidence verdict), in the Â§11.4.93 DB description column AND the docs AND the tracker, never a stub/Â§11.4.91-anti-pattern fragment; **(D3) BLOCKED items carry WHY + enumerated unblock CHOICES** “tightens Â§11.4.21: every Operator-blocked item (incl. Blocked/BLOCKED documented alias normalised to canonical Operator-blocked, never a silent fork Â§11.4.6) MUST enumerate the closed list of decisions/actions that would unblock it ([A] “ [B] “ [C] “, mirroring Â§11.4.66 “ “4-option shape); a BLOCKED item with no enumerated choices FAILs; **(D4) regular never-missed bidirectional DB “ docs “ tracker sync** “ the Â§11.4.93/95 git-tracked SQLite DB is the SSoT, docs + tracker DERIVED; full sync (db - to - md / md - to - db byte-identical round-trip Â§11.4.93 + tracker push) runs regularly + on every change, Â§11.4.86 drift-proof fingerprint (sha256 of the sorted item keyset, NOT mtime) gates freshness, Â§11.4.106 docs_chain-bound so drift is mechanically caught not vigilance-dependent; **(D5) generic external-tracker push** carries statuses (collapsed onto the tracker “ native set per Â§11.4.33/.112 when fixed, precise value preserved in a header, never lost), types, assignee (Â§11.4.104 participant handle, UNSET defaults from a project env var, never hardcoded/logged Â§11.4.10), and sub-tasks, IDEMPOTENT (dry-run-then-real, match-by-stable-key [<ID>] prefix / id custom-field, present “ UPDATE/absent “ CREATE, rate-limited, credential-redacted, sink-side created=N updated=M failed=0 proof Â§11.4.69), the MACHINERY project-agnostic Â§11.4.28 (consumer registers its tracker/list/field-map at runtime). Anti-bluff Â§11.4: every sync + validator pass carries captured evidence Â§11.4.5/.69; honest boundary Â§11.4.6 “ guarantees well-formed items + agreeing surfaces, NOT that the underlying work is correct (still crosses Â§11.4.108/.40/.123). Composes Â§11.4.15/.16/.21/.33/.34/.54/.66/.86/.91/.93/.95/.104/.106/.112/.123/.10/.28/.69/.6/Â§11.1. Classification: universal (Â§11.4.17) “ consumer supplies its DB path, id prefix, tracker service + list/board id + field map + default-assignee env var, docs_chain context per Â§11.4.35. Propagation gate CM-COVENANT-114-148-PROPAGATION (literal 11.4.148) + recommended gates CM-ITEM-INTEGRITY-STATUS-TYPE-ID / CM-ITEM-COMPREHENSIVE-DESCRIPTION / CM-BLOCKED-UNBLOCK-CHOICES / CM-TRACKER-SYNC-IDEMPOTENT + paired Â§1.1 mutation (gate-code = separate work item). **Canonical authority:** constitution submodule [Constitution.md](#) Â§11.4.148. Non-compliance is a release blocker. No escape hatch “ no --item-without-status, --item-without-type, --item-without-id, --stub-description-OK, --blocked-without-choices, --skip-tracker-sync, --one-way-sync-OK, --non-idempotent-tracker-push, --hardcode-assignee flag.

Â§11.4.149 “ Per-workable-item testing-diary mandate (User mandate, 2026-06-10).

Every workable item MUST carry an append-only TESTING DIARY “ a chronological record of every test run against it “ distinct from Â§11.4.93 item_history (lifecycle STATE transitions) + Â§11.4.55 Reopens (reopen cycles): the diary records TEST EXECUTIONS (which may or may not change status). The mandate (ALL hold): **(a)** a test_diary table additive to the Â§11.4.93/95 SQLite SSoT, one row per run soft-keyed to the item id, capturing ISO-8601-UTC date_time + tested_by (closed set User | Operator | AI-agent | HelixQA) + result (Â§11.4.45 PASS | FAIL | SKIP + detail) + in-depth observations (long markdown, facts or explicit UNCONFIRMED:/PENDING_FORENSICS: Â§11.4.6, the background a future fix needs) + action_taken+status_changed+from/to (did this run change status + WHY) + Â§11.4.69 evidence_path+feature_class, with a SCHEMA CONSTRAINT making a PASS row WITHOUT a non-empty evidence path impossible (a PASS-bluff rejected by the schema itself); **(b)** BOTH an in-depth diary doc + a derived at-a-glance summary VIEW (total/pass/fail/skip + last-verdict + last-run + status-changes + distinct testers/feature-classes, DERIVED never duplicated Â§11.4.93) feeding Â§11.4.132 risk-ordering; **(c)** four-format per-item Diary/Diary_Summary exports Â§11.4.65 + Â§11.4.86-drift-proof-

fingerprinted (sha256 of the sorted diary keyset) + Â§11.4.106 docs_chain-bound so a diary row not exported/pushed FAILs a gate (never-missed); (d) external-tracker SUB-TASK model â€” the item taskâ€™s description stays the item (Â§11.4.148 D2, NOT polluted with diary text), each run a CHILD sub-task (type task) with a {TODO|In-progress|Completed} lifecycle (collapsed per Â§11.4.33/.112), observations + action + an Evidence: line (path not raw artefact Â§11.4.10/.13) + a diary-entry idempotency key, reusing the Â§11.4.148 D5 idempotent rate-limited credential-redacted sink-side-proven push, mapper project-agnostic Â§11.4.28; (e) MINIMAL-LLM deterministic bash/Go tooling (zero LLM in the data path â€” the observations prose is authored by whoever ran the test; tooling only stores/renders/pushes/validates), 100% test-covered Â§11.4.27 (unit + integration-against-the-real-tracker with an honest Â§11.4.3 SKIP-when-token-absent never a faked PASS + export + HelixQA Challenge + paired Â§1.1 mutation) so a diary PASS-bluff is mechanically impossible. Honest boundary Â§11.4.6 â€” the diary guarantees a complete auditable per-item test-history, NOT that any single runâ€™s verdict is correct (rests on its own Â§11.4.69/.107/.123 evidence). Composes Â§11.4.6/.27/.45/.50/.55/.65/.69/.86/.93/.95/.106/.107/.123/.132/.148/.10/.13/.28/.33/.112/Â§1.1. Classification: universal (Â§11.4.17) â€” consumer supplies its DB path, diary doc layout, export formats, tracker sub-task field map, docs_chain context per Â§11.4.35. Propagation gate CM-COVENANT-114-149-PROPAGATION (literal 11.4.149) + recommended gates CM-TEST-DIARY-SYNC / CM-DIARY-PASS-REQUIRES-EVIDENCE + paired Â§1.1 mutation (gate-code = separate work item). **Canonical authority:** constitution submodule [Constitution.md](#) Â§11.4.149. Non-compliance is a release blocker. No escape hatch â€” no --skip-testing-diary, --diary-without-evidence, --no-diary-summary, --diary-without-tracker-subtask, --llm-in-diary-data-path, --diary-export-optional flag.

Â§11.4.150 â€” Mandatory deep multi-angle web research per change/issue, before declaring fixed or structural (User mandate, 2026-06-11). Verbatim operator mandate: â€œFor every single issue we fix or improvement we make â€” besides tight systematic-debugging, fixing, code review by independent agents, comprehensive tests â€” we MUST ALWAYS do deep web research from various angles! No matter how big/small/simple/complex, dig deep the internet for articles, technical documentation, APIs, open-source code â€” EVERYTHING that can help make the best possible solution OR confirm we donâ€™t have a more serious problem weâ€™re unaware of! We MUST do everything possible to STOP the constant issue-reopening and finally start closing items as fixed+working! Do this ALWAYS in parallel with the main work stream!â€” For EVERY fix / improvement / change / closure â€” no matter how big, small, simple, or complex â€” the agent MUST, IN ADDITION to tight systematic-debugging (Â§11.4.102), the fix, independent-agent code review (Â§11.4.125 / Â§11.4.142), and comprehensive multi-layer tests (Â§11.4.4(b) / Â§11.4.40), perform a DOCUMENTED **deep multi-angle web research pass** digging the internet from VARIOUS angles â€” articles, official + vendor technical documentation, API references, standards, issue trackers, maintainer guidance, reusable open-source code â€” to BOTH (i) discover the best possible solution AND (ii) confirm there is NOT a more serious underlying problem the team is unaware of. Forensic case study (FACT, 2026-06-11): a subtitle-on-secondary-display goal verdicted Won't-fix: structurally-impossible (Â§11.4.112 secure-surface pixel-blanking) was OVERTURNED by a deep multi-angle research pass that surfaced the real OCR-of-the-PRIMARY-display path â€” a â€œwe canâ€™tâ€” became a shipped capability; the canonical anti-pattern of a structural verdict reached for want of research. The mandate (ALL hold): (A) **No closure-as-fixed/structural WITHOUT a documented deep-research pass** â€” no item may be marked Fixed/Implemented/Completed (Â§11.4.33) OR classified structurally-impossible wonâ€™t-fix (Â§11.4.112) until a deep multi-angle pass is documented; Â§11.4.150 makes Â§11.4.8â€™s pass UNCONDITIONAL (every item, however trivial, at the closure AND structural-verdict gates). (B) **Multiple angles, not a single lookup** â€” â‰¥ 2 genuinely-distinct angles (official-docs / standards / known-bug-trackers / alternative-approach / failure-mode / security / performance / platform-constraint / open-source-precedent) so a single-

source confirmation-bias miss (Â§11.4.145) cannot pass; a one-link drive-by is NOT deep multi-angle. **(C) Confirm-no-bigger-problem, not just find-a-fix** â€” explicitly seek evidence the fix does not mask a deeper defect AND the closure/verdict is genuinely safe; â€œwe found nothing worseâ€” requires the enumerated search (Â§11.4.118). **(D) Latest-source + cited** â€” LATEST authoritative versions per Â§11.4.99 (never training-data/memory), each cited by URL + access date in the research artefact AND the closure commit footer (Deep-research <date>: <urls> OR the literal NO external solution found â€” original work per Â§11.4.8). **(E) Reopen-breaking is the PURPOSE** â€” STOP the constant Fixedâ†”i,ŽReopened churn (Â§11.4.34 / Â§11.4.55); a reopen whose root cause a deep pass would have surfaced is a Â§11.4.150 miss; composes Â§11.4.7 (demotion-evidence) + Â§11.4.112 (structural verdict now REQUIRES the cited-authorities pass). **(F) ALWAYS in parallel with the main stream** â€” background subagent-driven (Â§11.4.70 / Â§11.4.20 / Â§11.4.103 / Â§11.4.89) concurrent with main fix/build/test work, NEVER serialising it or stalling the loop (Â§11.4.94 / Â§11.4.97 / Â§11.4.101 / Â§11.4.126). **(G) Apply ASAP to every workable item** â€” retroactively + going forward across the Â§11.4.93 / Â§11.4.95 SSoT; items closed without a documented pass are re-audited in the Â§11.4.40 / Â§11.4.42 release-gate sweep. Honest boundary (Â§11.4.6): the pass reduces the unknown-unknown surface + breaks the most common reopen causes â€” it does NOT prove zero remaining defects (Â§11.4.118) and does NOT replace Â§11.4.108 runtime-signature verification, Â§11.4.125 / Â§11.4.142 review, or Â§11.4.40 retest; it is the research layer every fix/closure/structural-verdict additionally crosses, and â€œwe probably donâ€™t have a bigger problemâ€” without the enumerated multi-angle search is a guess (Â§11.4.6), never a finding. Composes Â§11.4.8 / Â§11.4.99 / Â§11.4.123 / Â§11.4.118 / Â§11.4.145 / Â§11.4.125 / Â§11.4.142 / Â§11.4.7 / Â§11.4.112 / Â§11.4.34 / Â§11.4.55 / Â§11.4.70 / Â§11.4.20 / Â§11.4.89 / Â§11.4.103 / Â§11.4.40 / Â§11.4.42 / Â§11.4.93 / Â§11.4.95 / Â§11.4.6 / Â§1.1. Classification: universal (Â§11.4.17) â€” the consuming project supplies its concrete research corpora, angle set, item tracker, and closure-commit-footer convention per Â§11.4.35. Propagation gate CM-COVENANT-114-150-PROPAGATION (literal 11.4.150) + recommended gate CM-DEEP-RESEARCH-PER-ISSUE (every closed/structural-verdicted item carries a documented multi-angle deep-research artefact + cited-source closure footer, run in parallel, before the closure/structural verdict is accepted) + paired Â§1.1 mutation (strip the literal â†” propagation gate FAILs; close an item or reach a structural verdict with no documented multi-angle research artefact / cited footer â†” CM-DEEP-RESEARCH-PER-ISSUE FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule [Constitution.md](#) Â§11.4.150. Non-compliance is a release blocker. No escape hatch â€” no --skip-deep-research, --single-source-suffices, --trivial-change-no-research, --close-without-research, --structural-without-research, --serialise-research, --research-from-memory-OK flag.

Â§11.4.151 â€” Project-prefixed release-tag/version-naming mandate (User mandate, 2026-06-12). Verbatim operator mandate: â€œEvery release tag and version name we create â€” on the main repository and on every Submodule we own â€” MUST be prefixed with the projectâ€™s release prefix, e.g.Â helix_translate-1.0.0-dev-0.0.1. The prefix MUST come from HELIX_RELEASE_PREFIX in our .env if it is set, otherwise from the lowercased project root directory name. The SAME prefix MUST be used across the main repo and all owned Submodules in one release so a release is greppable across every repository.â€” Every release tag AND every version name created on the main repository AND on every owned-by-us submodule (Â§11.4.28) MUST be prefixed with the projectâ€™s release prefix, form <PREFIX>-<version> (canonical example: helix_translate-1.0.0-dev-0.0.1), so a release is identifiable + greppable across every repository it spans (one git tag -l ' <PREFIX>-* ' enumerates the whole release surface). **Prefix resolution order (closed-set, deterministic â€” Â§11.4.6 no-guessing):** (1) HELIX_RELEASE_PREFIX from the projectâ€™s .env â€” authoritative when set; .env is git-ignored per Â§11.4.30 and the variable is documented in the tracked .env.example (a Â§11.4.77 re-obtain mechanism,

never committed); (2) fallback = the lowercased snake_case form of the project root directory name (no spaces) per [§11.4.29](#) used whenever the env var is unset/empty, so a prefix is ALWAYS resolvable from the checkout with zero operator input. **The prefix MUST be IDENTICAL across the main repo and all owned submodules within a single release** a release that tags the main repo `<PREFIX>-1.2.0` while tagging an owned submodule with a different/unprefixed value is a [§11.4.151](#) violation (the cross-repo grep no longer enumerates the release). Version codes increment monotonically within the prefix (`<PREFIX>-0.0.1` → `<PREFIX>-0.0.2`), never reset, never skipped. Honest boundary ([§11.4.6](#)): the prefix guarantees a release is identifiable + uniform across every repository, NOT that its contents are correct the tag is still created only after the [§11.4.40](#) full-suite retest GREEN and reaches every upstream via the [§11.4.113](#) merge-onto-latest-main path (NEVER a force-push), fanned out per [§2.1](#). Composes [§2.1](#) / [§11.4.29](#) / [§11.4.30](#) / [§11.4.40](#) / [§11.4.113](#) / [§11.4.126](#) (the release-scope terminal condition is a published, prefixed tag). Classification: universal ([§11.4.17](#)) the consuming project supplies its concrete prefix value + the `HELIX_RELEASE_PREFIX` env var per [§11.4.35](#). Propagation gate `CM-COVENANT-114-151-PROPAGATION` (literal `11.4.151`) + recommended gate `CM-RELEASE-PREFIX-NAMING` (every release tag/version on the main repo + every owned submodule carries the resolved `<PREFIX>-prefix`, identical across the release) + paired [§1.1](#) meta-test mutation (strip the literal `'` propagation gate FAILs; create an unprefixed or differing-prefix release tag `'` `CM-RELEASE-PREFIX-NAMING` FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule [Constitution.md](#) [§11.4.151](#). Non-compliance is a release blocker. No escape hatch no `--no-release-prefix`, `--unprefixed-tag`, `--prefix-optional`, `--differing-submodule-prefix` flag.

[§11.4.152](#) Crashlytics-recorded-data continuous monitoring + systematic-debug + regression-test-coverage mandate (User mandate, 2026-06-13). Verbatim operator intent: “For every project that has Firebase Crashlytics enabled / wired, we MUST continuously monitor ALL of the Crashlytics-recorded data crashes, ANRs, performance traces, and non-fatals systematically debug each, fix and improve, and cover everything with validation and verification tests. This MUST be checked regularly, with no false results and no bluff of any kind! ♦ Every project that has Firebase Crashlytics enabled/wired (SDK linked into a shipping artifact, crash+non-fatal+ANR reporting active) MUST treat the Crashlytics console as a first-class captured-evidence channel from real end-user devices and continuously process every datum it records an open Crashlytics issue with a green test suite is a [§11.4](#) PASS-bluff at the field-telemetry layer (a real user hitting a broken feature while the suite reports green). STRENGTHENS+COMPLETES [§11.4.47](#): [§11.4.47](#) owns the periodic REVIEW + dedup + Issue-creation pass that SURFACES Crashlytics/Analytics/Performance findings; [§11.4.152](#) owns what happens to each surfaced item AFTER the systematic-debug `'` fix/improve `'` regression-test-coverage lifecycle that drives it to a proven regression-immune closure ([§11.4.47](#) finds it; [§11.4.152](#) fixes it + proves it stays fixed; both mandatory, neither substitutes). The mandate (ALL hold): (1) **continuous monitoring of ALL four surfaces** fatal crashes, ANRs, performance traces/regressions, AND non-fatals (skipping any one is a PASS-bluff; non-fatals are the silent class every catch/fallback/recovered-error path is a real degraded UX and MUST be triaged, not ignored because the app did not crash); (2) **systematic-debugging of each (reproduce-before-fix)** per [§11.4.102](#) Iron Law + [§11.4.115](#) RED-baseline-on-the-broken-artifact (the recorded stacktrace/ANR-trace/non-fatal-context IS the [§11.4.5/69](#) captured defect-present evidence) BEFORE the fix a fix without a prior reproducing test is a [§11.4.43/123](#) violation; unreproducible `'` deep-research-before-declaring-untestable [§11.4.123/150](#); (3) **a fix/improvement for every confirmed issue** against its proven root cause ([§11.4.9/43/108](#)), acknowledged in console ♦ is not a fix, muting a recurring issue without a tracked rationale is a [§11.4.6/90](#) violation; (4) **validation+verification regression-test coverage per closed issue** in the SAME commit as the fix register a permanent [§11.4.135](#) regression guard ([§11.4.115](#) polarity test: `RED_MODE=1` captures the recorded defect on the pre-fix artifact, `RED_MODE=0` standing

GREEN guard asserting ABSENT) exercising the same user-reachable path the stacktrace identifies, with rock-solid captured evidence Â§11.4.5/.69/.107/.123 + a closure log citing the console issue id/URL + root-cause + fix commit + the validation/verification test paths; a Crashlytics issue marked resolved WITHOUT a falsifiable regression test is FORBIDDEN (the silent-recurrence vector, Â§11.4.138 operator-escape class applied to field telemetry); (5) **regular cadence** reusing the Â§11.4.47 five-trigger set (pre-build/pre-flash/pre-distribute/pre-tag blocking; daily + post-deployment-burn-in non-blocking) â€” a one-time sweep never repeated is a violation; (6) **no false results, no bluff** (Â§11.4.1/.6 â€” â€œno new issuesâ€” requires the enumerated monitored-surfaces+window Â§11.4.118; â€œfixedâ€” requires the REDâ†’GREEN flip Â§11.4.115/.130; a guard whose Â§1.1 mutation does not FAIL it is a bluff gate). Honest boundary Â§11.4.6: processing every recorded issue breaks the silent-recurrence vector â€” it does NOT prove zero remaining field defects nor replace Â§11.4.108/.40; an issue is â€œclosedâ€” only when its guard is GREEN on a clean artifact, never when the console mark is flipped. Classification: universal (Â§11.4.17) â€” the consuming project supplies its Crashlytics handle, console-access credential (Â§11.4.10, never logged), severity table, and per-issue closure-log path per Â§11.4.35; the reference project-level instantiation is the Lava clientâ€™s Â§6.O (Crashlytics-Resolved Issue Coverage â€” per-issue validation test + Challenge test + .lava-ci-evidence/crashlytics-resolved/<date>-<slug>.md closure log) + Â§6.AC (Comprehensive Non-Fatal Telemetry â€” every catch/fallback records a non-fatal with triage context). Composes Â§11.4.1/.5/.6/.9/.34/.40/.43/.47/.69/.90/.102/.107/.108/.115/.118/.123/.130/.135/.138/.150/ Â§1.1. Propagation gate CM-COVENANT-114-152-PROPAGATION (literal 11.4.152) + recommended gate CM-CRASHLYTICS-ISSUE-FULLY-COVERED (every closed Crashlytics issue carries a systematic-debug root-cause record + a registered Â§11.4.135 regression guard + a closure-log entry citing the console issue id/URL + the validation/verification test paths; a console-resolved issue with no falsifiable regression guard FAILs) + paired Â§1.1 meta-test mutation (gate-code = separate work item). **Canonical authority:** constitution submodule [Constitution.md](#) Â§11.4.152. Non-compliance is a release blocker. No escape hatch â€” no --skip-crashlytics-monitoring, --monitor-crashes-only, --skip-non-fatals, --resolve-without-regression-test, --console-mark-is-fixed, --monitor-once, --mute-without-rationale flag.

Â§11.4.153 â€” Comprehensive per-feature Status + Status_Summary document set with mandatory video-recording confirmation (User mandate, 2026-06-15). Every project MUST maintain, under docs/features/, a comprehensive feature Status document set (Status.md + its Â§11.4.56 Status_Summary.md companion) enumerating EVERY system component, EVERY client app/binary/surface (TUI / CLI / Web / desktop / mobile / API / gRPC / library / submodule / infrastructure), and EVERY feature â€” including features ported from any incorporated CLI-agent / submodule catalogue (Â§11.4.74) â€” with NO single feature left out. (1) **Total categorized coverage** â€” per-feature table organized Componentâ†’Category, reconciled against the actual codebase (a code-present feature missing from the table, or a row with no code, is a Â§11.4.6/Â§11.4.118 bluff). (2) **Per-feature fields** â€” Component / Feature / Category / Implementation / Wiring (genuinely end-user-reachable, not merely compiled, Â§11.4.108) / Real-use / Tests-coverage (four-layer Â§11.4.4(b)) / Validation (PASS / FAIL / SKIP / PENDING_FORENSICS / OPERATOR-BLOCKED per Â§11.4.45) / **Video-recording confirmation** (path to the real-use video or honest gap marker). (3) **Mandatory per-feature real-use video** â€” every user-visible â€œconfirmedâ€” claim backed by a recorded real-use video of a genuine end-user scenario (real prompts â†’ real LLM/service responses â†’ real results), NEVER a frozen/stale frame (Â§11.4.107), NEVER faked/mockd/demo-loop/bluff response or LLM error passed off as success (Â§11.4.2/Â§11.4.5), stored at the project-declared recording path (Â§11.4.35); a confirmed row with no real video or a bluff video = Â§11.4 PASS-bluff; autonomous-infeasible â†’ honest Â§11.4.3/Â§11.4.52 SKIP + tracked migration item, NEVER a faked confirmation. (4) **Video-analysis remediation loop** â€” every video analysed; any defect it surfaces triggers Â§11.4.102 systematic-debug â†’ fix â†’ Â§11.4.146 retest â†’ re-

record â†’ clean GO per Â§11.4.134 before â€œconfirmedâ€ (a video exposing a broken feature is a Â§11.4.4 test-interrupt). (5) **Always-in-sync** â€” Â§11.4.45-class roster/corpus-backed, Â§11.4.106 docs_chain-bound + Â§11.4.86 drift-proof fingerprint (sha256 of sorted feature-key roster AND sorted video-artefact roster, NOT mtime), re-syncs out-of-the-box; stale = violation. (6) **Four-format export** â€” HTML + PDF + **DOCX** (this doc class ADDS DOCX to the Â§11.4.65 HTML+PDF set; other classes unchanged), in sync per Â§11.4.60. (7) Follows Â§11.4.44/.45/.56/.57/.59/.60. (8) **MP4 format REQUIRED**. All video confirmations MUST be in .mp4 format (H.264). Window-specific capture ONLY (Â§11.4.159(A)). Vision validation REQUIRED (Â§11.4.159(D)). .cast files are supplementary only. Honest boundary Â§11.4.6 â€” guarantees a complete video-confirmed always-synced ledger, NOT per-feature correctness (rests on each rowâ€™s Â§11.4.69/.107/.123 evidence) and NOT a Â§11.4.40 retest substitute. Classification: universal (Â§11.4.17). Composes Â§11.4.2/.5/.44/.45/.52/.56/.57/.59/.60/.65/.86/.102/.106/.107/.108/.118/.123/.134/.146/Â§1.1. Propagation gate CM-COVENANT-114-153-PROPAGATION (literal 11.4.153) + recommended gates CM-FEATURE-STATUS-COMPLETE + CM-FEATURE-STATUS-VIDEO-CONFIRMED + paired Â§1.1 mutation.

Canonical authority: constitution submodule [Constitution.md](#) Â§11.4.153. Non-compliance is a release blocker. No escape hatch â€” no --skip-feature-status, --feature-without-video, --frozen-video-OK, --bluff-response-video-OK, --skip-video-analysis, --feature-ledger-incomplete-OK, --no-docx-export, --allow-feature-ledger-drift flag.

Â§11.4.154 â€” Window-scoped capture + fresh-corpus rotation for feature/QA recordings (User mandate, 2026-06-15). Verbatim: â€œrecording you perform does record the window containing apps and services, not the whole desktop or monitor screen!â€ + â€œAll old recording files MUST BE removed when new one starts!â€ Refines Â§11.4.2/.5/.107/.153 recording discipline with two capture-hygiene invariants. (A) **Window-scoped, NOT whole-screen** â€” every feature/QA video MUST capture ONLY the window/surface of the app/service under test (GUI window / CLI-TUI terminal pane / web tab-viewport / device-emulator-simulator frame), NEVER the whole desktop/monitor or unrelated windows; whole-desktop capture leaks operator-private content (Â§11.4.10/.83), dilutes the Â§11.4.107 liveness/freeze oracle, and breaks the Â§11.4.137 OCR/ROI content oracle. Target the window/region by stable identity (window id/title, device serial, browser context, tmux target) per Â§11.4.111 â€” never a fixed full-screen device index. Platform genuinely cannot capture below whole-screen â†’ honest Â§11.4.3 SKIP + tracked migration item, never a whole-screen pass-off. (B) **Fresh-corpus rotation** â€” when a new recording run for a scope begins, the agentâ€™s OWN prior in-scope stale recordings at the raw recording path MUST be removed FIRST so the live corpus reflects the current run (Â§11.4.107 not-stale + Â§11.4.86 roster-freshness). Honest boundary (Â§11.4.6 + Â§9.2): â€œremove oldâ€ = the agentâ€™s own prior recordings for the SAME scope/project ONLY â€” NEVER another projectâ€™s/scopeâ€™s/operator-authored files; uncertain â†’ surface, donâ€™t delete (Â§11.4.122); committed docs/qa/<run-id>/evidence (Â§11.4.83) is the durable record, NOT rotated. Classification: universal (Â§11.4.17). Composes Â§11.4.2/.5/.10/.83/.86/.107/.111/.122/.128/.137/.153/Â§9.2/.6. Propagation gate CM-COVENANT-114-154-PROPAGATION (literal 11.4.154) + recommended gate CM-WINDOW-SCOPED-FRESH-CORPUS-RECORDING + paired Â§1.1 mutation.

Canonical authority: constitution submodule [Constitution.md](#) Â§11.4.154. Non-compliance is a release blocker. No escape hatch â€” no --whole-screen-capture-OK, --skip-window-scope, --keep-stale-recordings, --no-corpus-rotation, --full-desktop-recording flag. (C) **MP4 auto-conversion REQUIRED**. Any .cast file produced MUST be auto-converted to .mp4 via agg + ffmpeg immediately after capture. The .mp4 is the primary evidence; .cast is supplementary only.

Â§11.4.155 "Project-name-prefixed feature/QA recording filenames (User mandate, 2026-06-15). Verbatim: "All recorded videos MUST START with prefix: the PROJECT NAME (ALWAYS USE THE PROJECT NAME). Project name MUST be obtained according to the constitution's own project-name resolution." Every recorded video the project produces "every feature/QA real-use recording (Â§11.4.153), every window-scoped capture (Â§11.4.154), every always-on device recording (Â§11.4.128), every raw or curated artefact at the project-declared recording path (Â§11.4.35) + the committed docs/qa/<run-id>/trail (Â§11.4.83) "MUST have a filename that STARTS WITH the PROJECT-NAME prefix, ALWAYS; an unprefixed recording is a Â§11.4.155 violation (a multi-project/scope corpus on one host per Â§11.4.128/.103 becomes un-greppable + un-attributable "the Â§11.4.151 identify-and-grep failure on the recording-corpus axis). **Prefix resolution (closed-set, deterministic** "Â§11.4.6, IDENTICAL to Â§11.4.151): (1) HELIX_RELEASE_PREFIX from .env (authoritative, git-ignored Â§11.4.30, documented in tracked .env.example Â§11.4.77) else (2) lowercased snake_case project-root dir name Â§11.4.29 "ALWAYS resolvable, zero operator input. SAME prefix for EVERY recording in a checkout so ls '<PREFIX>---'* enumerates the corpus; canonical form <PREFIX>---<feature-or-scope>---<run-id>.<ext> (--- delimits the prefix unambiguously); MUST equal the Â§11.4.151-resolved release-tag prefix (divergence is itself a Â§11.4.155 violation "one project, one name). Honest boundary (Â§11.4.6): the prefix guarantees attribution + greppability, NOT content validity (still Â§11.4.107/.137/.153) and does NOT relax Â§11.4.154's window-scope/rotation (rotation removes the agent's OWN <PREFIX>---* only; foreign/operator files surfaced not deleted Â§11.4.122/Â§9.2). Classification: universal (Â§11.4.17). Composes Â§11.4.151/.128/.153/.154/.111/.83/.6/.29/.30/.35/.77/.86/Â§1.1. Propagation gate CM-COVENANT-114-155-PROPAGATION (literal 11.4.155) + recommended gate CM-RECORDING-PROJECT-NAME-PREFIX + paired Â§1.1 mutation.

Canonical authority: constitution submodule [Constitution.md](#) Â§11.4.155. Non-compliance is a release blocker. No escape hatch "no --no-recording-prefix, --recording-without-project-name, --unprefixed-recording, --prefix-optional-for-recording, --differing-recording-prefix flag.

Â§11.4.156 "All CI/CD automation (GitHub Actions / GitLab pipelines / equivalents) MUST be disabled (User mandate, 2026-06-15). Verbatim operator mandate: "Any GitHub actions or GitLab pipelines MUST BE disabled! Add this critical mandatory rule / mandatory constraint into the root constitution Submodule, commit and fetch all its changes to all upstreams and make sure we respect and follow this rule (we do apply it) ASAP!!!" Every repository this Constitution governs "main repo, this constitution submodule, every owned + nested submodule we author and push "MUST ship with ALL server-side CI/CD automation DISABLED: no push to any owned upstream may trigger a GitHub Actions run, GitLab pipeline, or equivalent (Jenkins/CircleCI/Travis/Drone/Woodpecker/Bitbucket/Azure, any on: push/schedule/workflow_dispatch). GENERALISES + makes ABSOLUTE the Â§11.4.75 Layer-5 posture (remote CI DISABLED, workflow preserved at a "disabled-local-only non-.yaml name a provider ignores) across ALL governed repos; enforcement migrates to the LOCAL Â§11.4.75 git-hook ritual + Â§11.4.40 pre-tag sweep, never a remote runner. ALL hold: (A) zero active .github/workflows/*.yaml|yaml / .gitlab-ci.yml / .gitlab/**/ equivalent at the ROOT of any governed repo/submodule (the only place a provider executes); (B) "disabled" = a push triggers ZERO runs "delete OR rename to a non-trigger name (Â§11.4.75 .disabled/.disabled-local-only); a live-on: + if: false workflow still queues runs, NOT compliant; (C) scope = repos we author+push "vendored/third-party nested configs below the root (AOSP external/**, prebuilts/**, vendored submodules) are INERT (a provider never runs a non-root config), OUT of scope, MUST NOT be mass-edited (Â§11.4.29 vendor-exempt); test = "does a push to OUR upstream trigger a run?" yes? "disable, inert? "document+leave (Â§11.4.6 verify-not-assume); (D) no new CI may be added (release blocker); (E) pre-push verify git ls-files

| grep -E '^\.github/workflows/.*\ya?ml\$|^\.gitlab-ci\.yaml\$'
empty for authored repos, Â§11.4.109-class PreToolUse guard + gate enforce mechanically.
Honest boundary (Â§11.4.6): file-level disabling stops FILE-triggered runs, NOT provider-side
server settings (org-default required workflows, branch-protection required checks, provider
scheduled exports) â€” the operator turns those off; the agent documents what it cannot reach,
never claims unachieved completeness. Composes Â§11.4.75/.29/.6/.40/.42/.109/.113/Â§2.1.
Classification: universal (Â§11.4.17). Propagation gate CM-COVENANT-114-156-
PROPAGATION (literal 11.4.156) + recommended gate CM-NO-ACTIVE-CI + paired Â§1.1
meta-test mutation (strip the literal â†’ propagation gate FAILs; add a root .github/
workflows/x.yml to an authored repo â†’ CM-NO-ACTIVE-CI FAILs). **Canonical
authority:** constitution submodule [Constitution.md](#) Â§11.4.156. Non-compliance is a
release blocker. No escape hatch â€” no --allow-ci, --enable-workflow, --keep-
pipeline, --remote-ci-OK, --ci-exempt flag.

**Â§11.4.157 â€” GEMINI.md maintained in lockstep with CLAUDE.md / AGENTS.md /
QWEN.md (User mandate, 2026-06-15).** Verbatim operator mandate: â€œMake sure with
CLAUDE.md, AGENTS.md, QWEN.md we maintain GEMINI.md too! Add this mandatory
fact / rule to root constitution Submodule we are inheriting / extending - CONSTITUTION.md,
CLAUDE.md, QWEN.md, AGENTS.md, GEMINI.md and other related relevant files!â€
Forensic FACT (2026-06-15): when Â§11.4.156/Â§11.4.157 were authored, Constitution/
CLAUDE/AGENTS/QWEN carried the family through Â§11.4.155 but GEMINI.md had silently
drifted to Â§11.4.141 â€” 14 mandates (Â§11.4.142â€”155) never propagated â€” a Â§11.4
propagation-bluff (a Gemini-CLI agent reads a stale Constitution). GEMINI.md is a FIRST-
CLASS governance context carrier EQUAL to CLAUDE.md/AGENTS.md/QWEN.md, never
optional/best-effort. ALL hold: **(A)** five-carrier lockstep â€” no governance change is complete
until GEMINI.md carries it alongside the other three mirrors; GEMINI.md is added to the
Â§11.4.26 propagation + cross-reference set explicitly; **(B)** no silent drift â€” GEMINI.md
lagging the other mirrorsâ€™ highest rule is a Â§11.4.157 violation (Â§11.4.65-class), back-fill
required; **(C)** equal status â€” GEMINI.md restates the SAME literal 11.4.N anchors the
propagation gates require (Â§11.4.35), fleet count INCLUDES GEMINI.md; **(D)** consumer
projectsâ€™ own CLAUDE/AGENTS/QWEN/GEMINI bind too (Â§11.4.35). Honest boundary
(Â§11.4.6): the Â§11.4.142â€”155 GEMINI.md back-fill is a tracked release-blocking
remediation; claiming GEMINI.md â€œin syncâ€ while the back-fill is incomplete is itself a
Â§11.4.157 violation. Composes Â§11.4.26/.35/.17/.44/.65/.140/.156/Â§1.1. Classification:
universal (Â§11.4.17). Propagation gate CM-COVENANT-114-157-PROPAGATION (literal
11.4.157, GEMINI.md INCLUDED) + recommended gate CM-GEMINI-MD-LOCKSTEP +
paired Â§1.1 meta-test mutation. **Canonical authority:** constitution submodule
[Constitution.md](#) Â§11.4.157. Non-compliance is a release blocker. No escape hatch â€” no
--skip-gemini-md, --gemini-optional, --gemini-lag-OK, --four-
carrier-suffices flag.

**Â§11.4.158 â€” Intensive all-feature/flow/edge-case video-recording + read-the-screen
content-verification mandate (User mandate, 2026-06-16).** Every project MUST be covered by
intensive automated testing that exercises + RECORDS every feature/flow/use-case/edge-case
(valid/invalid/boundary/concurrent/failure-injection Â§11.4.85), each recording showing the
feature GENUINELY WORKING with REAL results (real promptsâ†’real responsesâ†’real
outputs Â§11.4.153) and NO false/simulated/stale/frozen result (Â§11.4.2/.5/.107) â€” a
recording showing an error/bluff/non-working feature is a FINDING (â†’Â§11.4.153(4)/Â§11.4.4
fixâ†’retestâ†’re-record), never a confirmation. The testing System MUST ACTUALLY READ
every shown log line / message / UI label / dialog / toast / status text and VERIFY it is a genuine
working result (OCR/ROI Â§11.4.117/.137 + confidence floor for pixel surfaces; direct capture
for terminal/log; Â§11.4.107(10) self-validated golden-good/golden-bad analyzer) â€” â€œa
video was producedâ€ is NOT evidence, â€œthe System read the screen + confirmed a genuine
resultâ€ is. HelixQA (Â§11.4.27) MUST drive this exerciseâ†’recordâ†’readâ†’score pass,

PASS only on a read-confirmed genuine result + captured artefact path (Â§11.4.69). Default recording save path = \$HOME/Downloads (host user's home Downloads, resolved at runtime never hardcoded) unless a project declares an override per Â§11.4.35; Â§11.4.155 project-prefix + Â§11.4.154 window-scope/rotation + Â§11.4.128 git-ignored raw corpus + Â§11.4.83 curated docs/qa evidence apply. **Vision analysis MANDATORY for every recording** after every recording, the agent MUST read the terminal output / video content and verify the feature actually works. A recording without a vision-confirmed verdict is a Â§11.4.158 violation. Composes Â§11.4.2/.5/.25/.27/.52/.69/.83/.85/.107/.108/.117/.118/.128/.137/.138/.153/.154/.155/Â§1.1. Classification: universal (Â§11.4.17). Propagation gate CM-COVENANT-114-158-PROPAGATION (literal 11.4.158) + recommended gates CM-INTENSIVE-RECORDING-COVERAGE + CM-RECORDING-CONTENT-READ-VERIFIED + paired Â§1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) Â§11.4.158. Non-compliance is a release blocker. No escape hatch "no --skip-recording-coverage, --video-without-content-read, --happy-path-recording-suffices, --recording-path-anywhere, --unread-recording-OK, --skip-helixqa-read flag.

Â§11.4.159 "Mandatory window-specific video recording + vision validation mandate (User mandate, 2026-06-20). Every feature test, validation, verification, challenge, and QA session that produces video evidence MUST comply with ALL of the following: **(A) Window-specific recording ONLY** "every video MUST record ONLY the target application window (Terminal pane, TUI app, browser tab, emulator frame), NEVER the whole desktop/monitor screen; use macOS screencapture -l<window_id>, Linux xdotool+ffmpeg, or equivalent; whole-desktop capture leaks operator-private content (Â§11.4.10), dilutes the Â§11.4.107 liveness oracle, and breaks Â§11.4.137 OCR/ROI; platform genuinely cannot capture below whole-screen => honest Â§11.4.3 SKIP + tracked migration item. **(B) MP4 format REQUIRED** ".mp4 (H.264, movflags +faststart, pix_fmt yuv420p); .cast files are supplementary only; auto-convert via agg+ffmpeg per Â§11.4.154(C). **(C) Project-name prefix REQUIRED** "filename MUST start with project name in snake_case from HELIX_RELEASE_PREFIX in .env (Â§11.4.151) or lowercased project root dir name (Â§11.4.29); format <project_name>-<feature>-<timestamp>.mp4; unprefixed = violation. **(D) Mandatory vision validation** "after EVERY recording, the agent MUST read the terminal output / video content and verify: (i) feature ACTUALLY WORKS, (ii) LLM responses are REAL, (iii) all tests show PASS, (iv) no "TODO implement" / "simulate" / "for now" patterns, (v) output demonstrates end-user working feature; MUST produce a verdict (PASS/FAIL) with evidence path (Â§11.4.69). **(E) Terminal window cleanup** "after each recording, dismiss/close ONLY the Terminal window used for that recording (window-specific close via osascript/xdotool); MUST NOT close windows belonging to other processes. **(F) Real results ONLY** "every recording MUST show REAL working features; errors/empty output/simulated responses trigger fix+retest+re-record. **(G) Re-runnable evidence** "the command shown MUST be re-runnable to produce the same results. **(H) Fresh-corpus rotation** "remove agent's own prior in-scope stale recordings FIRST (Â§11.4.154). **(I) Content verification MANDATORY** "not duration-based "the value of a recording is NOT its duration but its CONTENT; a recording MUST demonstrate the ACTUAL feature being used with REAL results; before accepting ANY recording, the agent MUST verify: expected output patterns ARE present (e.g., test PASS lines, API response data, feature-specific output), the feature ACTUALLY WORKS as demonstrated (not just "something ran", LLM responses are REAL content (not simulated, not placeholder, not empty), every claim of "working" is backed by visible evidence; a 5-second recording showing a feature working correctly is MORE valuable than a 60-second recording of empty terminal; duration is NOT a proxy for quality. **(J) Expected-content specification REQUIRED** "before recording, the agent MUST specify what content SHOULD appear (expected patterns, expected test results, expected API responses); after recording, the agent MUST verify these patterns ARE present; if expected content is MISSING, the recording is REJECTED regardless of

duration. **(K) Content-verification recording workflow** – mandated workflow: (1) SPECIFY expected content patterns, (2) RECORD the feature execution, (3) EXTRACT all text from the recording, (4) VERIFY expected patterns are present, (5) CHECK for simulated/placeholder content, (6) ACCEPT only if ALL patterns found AND zero bluffs detected, (7) REJECT and re-record if ANY pattern missing or bluff detected. **(L) Root cause analysis REQUIRED for rejected recordings** – when a recording is rejected (missing expected content, bluff detected, or empty capture), the agent MUST investigate WHY before re-recording per Â§11.4.102; determine the root cause (timing issue, wrong command, tool failure, etc.) and fix it; simply re-recording without understanding WHY the first attempt failed is a Â§11.4.159 violation. **(M) Real-time monitoring RECOMMENDED** – for complex features, use real-time monitoring that analyzes output DURING recording (not after); this catches issues immediately and allows corrective action before the recording completes. Classification: universal (Â§11.4.17). Composes Â§11.4.2/.3/.5/.10/.29/.69/.83/.107/.111/.128/.137/.151/.153/.154/.155/.158/Â§1.1. Propagation gate CM-COVENANT-114-159-PROPAGATION (literal 11.4.159) + recommended gate CM-WINDOW-VIDEO-VALIDATED + paired Â§1.1 mutation. **Canonical authority:** constitution submodule [Constitution.md](#) Â§11.4.159. Non-compliance is a release blocker. No escape hatch – no --whole-screen-ok, --cast-only, --skip-vision-validation, --no-cleanup, --simulated-recording-ok, --unprefixed-recording flag.

Â§11.4.160 – Vision-verified recording + HelixQA bridge mandate (User mandate, 2026-06-21). Compact summary: every video recording for feature/QA evidence MUST be processed through a vision/OCR pipeline that reads on-screen content and confirms expected results BEFORE acceptance; the recording system MUST provide a bridge feeding captured frames to HelixQA’s test infrastructure (or equivalent) for automated read-the-screen verification against SPECIFY-phase expected patterns (Â§11.4.159(J)); capture frames at 5s intervals, run OCR/vision with self-validated analyzer (Â§11.4.107(10)), compare text against patterns, produce per-frame PASS/FAIL with evidence path, surface failures immediately for re-record (Â§11.4.159(L)). Project-configured parameters per Â§11.4.35. Honest boundary: does NOT replace Â§11.4.5 quality analysis nor Â§11.4.108 runtime-signature; FLAG_SECURE uses Â§11.4.117 proxy oracle. Classification: universal (Â§11.4.17). Propagation gate CM-COVENANT-114-160-PROPAGATION + recommended gate CM-VISION-VERIFIED-RECORDING-BRIDGE + paired Â§1.1 mutation.

Â§11.4.161 – Rootless container runtime mandate (User mandate, 2026-06-21). Compact summary: every project MUST use Podman rootless (or equivalent) for ALL containerized workloads – rootful Docker/sudo FORBIDDEN unless Â§11.4.112 platform constraint documented; vasic-digital/containers Submodule (Â§11.4.76) sole orchestration layer via pkg/boot/pkg/compose/pkg/health; extend upstream per Â§11.4.74 for missing capabilities; integration tests boot infra on-demand. Honest boundary: does NOT replace Â§11.4.10 credentials nor Â§12.3 hygiene. Classification: universal (Â§11.4.17). Propagation gate CM-COVENANT-114-161-PROPAGATION + recommended gate CM-ROOTLESS-CONTAINER-RUNTIME + paired Â§1.1 mutation.

Â§11.4.162 – OpenDesign UI design system mandate (User mandate, 2026-06-21). Compact summary: every project with user-facing interfaces MUST use OpenDesign (<https://github.com/nexu-io/open-design>) as mandatory UI design system – NOT ad-hoc CSS; install as dependency, use token system for all color (light+dark from brand assets), typography, spacing, component tokens; extend upstream (Â§11.4.74) for unsupported patterns; every UI component ships light+dark + token-diff regression tests via visual regression; no element overlap/font collision/label overlay. Honest boundary: does NOT replace Â§11.4.27 functional testing nor Â§11.4.48/49 UI testing. Classification: universal (Â§11.4.17). Propagation gate CM-COVENANT-114-162-PROPAGATION + recommended gate CM-OPENDSIGN-UI-SYSTEM + paired Â§1.1 mutation.

Â§11.4.163 â€” Universal Media Validation (User mandate). Every recorded artifact MUST pass MEDIA VALIDATION pipeline â€” OCR/transcription/text parsing, compare vs SPECIFY-phase patterns (Â§11.4.159(J)), self-validated analyzer (Â§11.4.107(10)), structured verdict, pinpoint on FAIL, REAL-TIME trigger option. Paired Â§1.1 mutation on golden-bad fixture. Classification: universal. Propagation gate CM-COVENANT-114-163-PROPAGATION + paired Â§1.1 mutation.

Â§11.4.164 â€” Universal Auto-Propagation Hook (User mandate). Every constitutional pull triggers `post_update_hook.sh` â€” detects changed files, registers skills/MCP/hooks/scripts, logs summary. Consumer invokes after every pull. Classification: universal. Propagation gate CM-COVENANT-114-164-PROPAGATION + paired Â§1.1 mutation.

Â§11.4.165 â€” Universal Independent Verification Agent (User mandate). Every code/media/docs/config output passes independent verifier (Â§11.4.70/.20), zero-finding GO (Â§11.4.134). CODE: review+build+mutations+runtime-signature. MEDIA: pipeline+genuine+format. DOCS: exports+revision. CONFIG: syntax+schema+leaks. Self-validated Â§1.1 mutation. Classification: universal. Propagation gate CM-COVENANT-114-165-PROPAGATION + paired Â§1.1 mutation.

Â§11.4.166 â€” Universal Semgrep static analysis (User mandate). Verbatim: â€œsemgrep MUST be available and used by all projects governed by this Constitution for static analysis / security scanning of all code changes, before any commit or push to any upstream! semgrep MUST be installed on every development host, auto-configured, and available on the system PATH for ALL users.â€” Every consuming project MUST ensure semgrep is installed + configured: (1) `pip install --user semgrep` or via `scripts/semgrep/semgrep_setup.sh`; (2) PATH via `scripts/semgrep/semgrep_path.sh` sourced from `.bashrc/.zshrc`; (3) scan before commit+push via `semgrep scan --config auto --error` â€” blocks on findings; (4) MCP integration per Â§11.4.78 step 3; (5) anti-bluff validation via `scripts/semgrep/semgrep_validate.sh`; (6) docs_chain context; (7) inherited by reference per Â§11.4.28/Â§11.4.35. Classification: universal. Propagation gate CM-COVENANT-114-166-PROPAGATION + paired Â§1.1 mutation.

When in doubt

- Read `constitution/Constitution.md` for the canonical text of every rule.
- Cross-reference the sibling carriers (`CLAUDE.md` / `AGENTS.md` / `QWEN.md`) â€” they carry the same universal content for their respective agents.
- If still unclear, ask the operator. Do NOT guess. Do NOT bluff (Â§11.4.6).