GEMINI

Helix Constitution — Universal GEMINI.md

Field	Value
Revision	25
Created	2026-06-09
Last modified	2026-07-17T05:00:00Z
Status	active

Status summary	Strengthened §11.4.115 + §11.4.135 + §11.4.146 (research-derived, 2026-07-17 — the status-without-custody forensics from a consuming project): NO new anchor minted — the rules already existed as prose and did not bind; landed the MECHANICAL BINDING instead. Forensic FACTs (genericised): 576 tracked items, 182 done-claiming, 173 (95%) with NO registered guard — an operator sampled 6 done-claimers, 6/6 broken (the expected draw of that arithmetic, not bad luck); recorded reopen-per-recorded-fix 52% (elite <10%), itself an UNDERCOUNT because recurrences re-entered as new ids and statuses moved with zero history rows; the guard suite's exit code read ONLY its FAIL count — measured live on a release candidate: 61 guards → PASS 0 / FAIL 5 / PENDING 56 → exit 0 "not blocked" — and the release-tag tooling consulted the guards ZERO times, so "never verified" and "verified good" were the same colour; the project's own history selected the discriminator: every fix whose done-claim was preceded by an on-device RED→GREEN flip that HELD recorded zero reopens, every fix confirmed at source-green bounced. §11.4.115(F) — machine-written verdict pairs + guard-viability: RED/GREEN are harness-written verdict files (target fingerprint read from the target at run time, RED fingerprint ≠ GREEN fingerprint, iterations ≥3, evidence class matches the defect's layer — a grep transcript can never satisfy a user-visible class); a guard never observed FAILing on the genuinely-broken artifact is unvalidated instrumentation and mints no verdicts; the canonical §1.1 mutation is the fix-commit's revert, and a string-deletion mutation that only removes what the gate greps is a refused tautology; gate `CM-GUARD-VERDICT-MACHINE-WRITTEN`. §11.4.135 verdict-coverage extension — the suite exit MUST incorporate COVERAGE, never only N_FAIL; the release seam blocks on `uncovered = registered ∧ topology-present ∧ no-verdict-for-the-candidate-fingerprint` (ABSENCE of a verdict blocks exactly as a FAIL does); the release-tag tooling MUST consult the verdict store; REQUIRES-REBUILD PENDING is self-clearing on the candidate, topology exemptions enumerated via a checked-in map never silent, adoption via a monotone-decrease ratchet, and a naive "PENDING always blocks" is rejected as a §11.4.201 FAIL-bluff; gate `CM-RELEASE-VERDICT-COVERAGE`. §11.4.146(D3) — status-custody: a done/ready status write is REFUSED at the status-write seam without the file chain registry-row-keyed-by-the-item's-EXACT-id (alias map for legacy keys) → executable guard → RED+GREEN verdict pair → class-matched evidence (prose nowhere in the chain); a full-table sweep gate (not a diff check) catches raw-DB by-

passes; §11.4.205 construction order — the executable sweep + its golden-bad fixture land in the SAME commit as any governance text citing them, and the meta-test asserts the sweep is WIRED; gate `CM-STATUS-CUSTODY`. Honest boundary §11.4.6 carried explicitly: the binding does not prove features bug-free (§11.4.186 family dedup remains), does not fix a miscalibrated oracle (§11.4.107(10) goldens necessary-not-sufficient), does not replace §11.4.185 manual QA, and makes target availability a hard dependency BY DESIGN. All Classification: universal (§11.4.17) — no hardware/vendor/project literal; consumers supply registry/sweep/topology-map paths as DATA per §11.4.35. Propagation gates for literals `11.4.115` / `11.4.135` / `11.4.146` unchanged (stay GREEN). CLAUDE.md + AGENTS.md + QWEN.md + GEMINI.md mirrors updated in lockstep (§11.4.157). Prior round: Added §11.4.213 — FEATURE research-scheduling directive (operator-directed mandate, 2026-07-16): a new §11.4.140-registered action `FEATURE` (all six grammar forms incl. the single-colon `FEATURE:` form) that SCHEDULES — never synchronously runs — a deep, enterprise-grade research + implementation-planning effort as a Type=Task workable item, reusing the §11.4.202 `report_item.sh` engine for item-creation + DB-sync + tracker-push (never reimplemented) plus a new durable `docs/requests/feature_queue.md` queue so a scheduled request is never dropped; the created item embeds the full mandate (deep web research §11.4.8/.99/.150; systematic-debug of every weak-spot/gap/danger-zone §11.4.102 with a risk-free rock-solid design per gap; always enterprise-grade/bleeding-edge/innovative; exhaustive documentation down to lines-of-code + micro-POCs + diagrams + SQL + templates §11.4.65/.73; reuse-first investigation of vasic-digital/HelixDevelopment components before rewriting, kept decoupled §11.4.28/.74/.177; full test-type + Challenges + HelixQA planning §11.4.27/.169; fine-grained phased/task/subtask breakdown; enterprise-scalability + max-performance planning; OpenDesign UI/UX wireframes where applicable §11.4.162/.190; full CodeGraph integration §11.4.78/.79/.80; fully-detailed workable items synced in real time to the SQLite SSoT + every connected external tracker §11.4.93/.95/.148 D5, honest-skipping any absent tracker §11.4.10) as its acceptance-defining work program — the multi-day research itself is EXECUTED by the standing autonomous loop (§11.4.87/.94/.97/.103/.126) when it claims the item, driven to a genuinely COMPLETED-and-wired or explicitly evidence-backed CLOSED terminal state under the §11.4.197 research-completion mandate, NEVER left sitting un-wired in the backlog. Propagation gate `CM-COVENANT-114-213-PROPAGATION` (literal

Field	Value
	`11.4.213`) + recommended gate `CM-FEATURE-DIRECTIVE` + paired §1.1 mutation. Classification: universal (§11.4.17). CLAUDE.md + AGENTS.md + QWEN.md + GEMINI.md mirrors updated in lockstep (§11.4.157). Prior round: Added §11.4.211 + §11.4.212 (User mandate 2026-07-16). §11.4.211 — Merge-conflict resolution during any main→feature/product/flavor merge MUST run on the Fable model at xhigh effort (Opus xhigh fallback): the merge-CONFLICT-resolution analogue of §11.4.209's code-REVIEW model-pin, generalised per §11.4.17 to EVERY track's controlled main→(feature
Issues	none
Issues summary	—
Fixed	§11.4.140 + §11.4.141 mirrors
Fixed summary	Carrier created in lockstep with the Constitution.md §11.4.140 + §11.4.141 additions.
Continuation	—

How inheritance works

This GEMINI.md is the **Gemini CLI carrier** of the Helix Constitution submodule. It is one of the four canonical agent context carriers (`constitution/CLAUDE.md` , `constitution/AGENTS.md` , `constitution/QWEN.md` , `constitution/GEMINI.md`) — these ARE the canonical root per §11.4.35, NOT consumer extensions. The full universal rule set lives in `constitution/Constitution.md` ; the sibling carriers (`CLAUDE.md` / `AGENTS.md` / `QWEN.md`) carry the same content for their respective agents. A consuming project's own repo-root `GEMINI.md` (if any) opens with the inheritance pointer and adds only project-specific rules; THIS file is the canonical carrier and does not.

For agents driven from a consuming project, the project's repo-root carrier SHOULD start with a clearly-marked inheritance pointer:

```
## INHERITED FROM constitution/GEMINI.md

All rules in `constitution/GEMINI.md` (and the `constitution/Constitution.md`
it references) apply unconditionally. The project-specific rules below
extend them.
   1 1    4    1 4 0                    1
```

Action-prefix recognition (§ . . — LAYER-)

The block below is the **canonical LAYER-1 action-prefix recognition instruction**, embedded VERBATIM (identical text) into every agent context carrier the constitution maintains. That textual identity across all four carriers IS the LAYER-1 universality: Gemini CLI reads this block via this carrier and applies the prefix expansion itself, so the action-prefix system works on Gemini CLI with zero extra setup.

§11.4.140 — Universal action-prefix system (`ACTION_NAME ::`) (User mandate, 2026-06-09; GRAMMAR_ADDENDUM 2026-06-09; arrow form 2026-07-02). When a user prompt's FIRST non-blank line starts with a recognised action prefix, you MUST: (1) look the action token up in the action registry `constitution/actions/registry.yaml` (or `$HELIX_ACTION_REGISTRY`); (2) if it is a registered action, REPLACE the prefix with that action's `expansion` text and apply its `rules` ; (3) execute the remainder of the prompt under the expanded instruction. **Five EQUIVALENT forms** — same action, same expansion, same execution: (1) `ACTION_NAME :: <rest>` (bare `::`), (2) `PREFIX::ACTION_NAME :: <rest>` (namespaced `::`), (3) `/ACTION_NAME <rest>` (bare slash), (4) `/PREFIX::ACTION_NAME <rest>` (namespaced slash), (5) `ACTION_NAME ---> <rest>` (bare arrow) ≡ `PREFIX::ACTION_NAME ---> <rest>` (namespaced arrow). Thus `BACKGROUND :: x` ≡ `DEFAULT::BACKGROUND :: x` ≡ `/BACKGROUND x` ≡ `/DEFAULT::BACKGROUND x` ≡ `BACKGROUND ---> x` ≡ `DEFAULT::BACKGROUND ---> x` . `PREFIX` is an action NAMESPACE; the reserved default namespace is **`DEFAULT`** , and an action runs WITH or WITHOUT the prefix. Grammar (all hold): anchored at the FIRST non-blank line only (mid-prose tokens never match); the action token AND the namespace are UPPERCASE-only `[A-Z][A-Z0-9_]*` (lowercase never matches); the namespace separator `::` inside the token carries NO surrounding spaces (`PREFIX::ACTION_NAME`), DISTINCT from the action-body separator `" :: "` (one ASCII space on each side of `::` — avoids C++ `Foo::Bar` , YAML `key: value` , URLs) in forms 1/2, the arrow-body separator `" ---> "` (one ASCII space on each side) in form 5, and the slash-body separator (one space) in forms 3/4; stacked prefixes (`A :: B :: rest`) apply outer-to-inner, left-to-right (expand `A` , re-scan, expand `B` , then the residual is the task); a leading `\` escapes the prefix for the `::` , arrow, and slash forms (`\BACKGROUND :: x` , `\BACKGROUND ---> x` , `\/BACKGROUND x` — treat literally, strip the backslash, NO expansion) so action names can be discussed. **Conflict rule (slash form):** `/ACTION_NAME` (form 3) is honored as the action ONLY when `ACTION_NAME` (case-folded) does not collide with a built-in/host slash command (registry `slash_bare: auto` + `slash_conflicts: [..]`); form 4 (`/PREFIX::ACTION_NAME`) is ALWAYS unambiguous and always honored. An unknown token that matches the grammar shape (any of the 5 forms) but is NOT registered is NEVER silently expanded or silently dropped — ask which registered action was meant (§11.4.66 / §11.4.105) or treat it as a literal prompt, NEVER invent an expansion (§11.4.6); any prompt not satisfying the grammar is an ordinary prompt and the system is a no-op. The registered action **`BACKGROUND`** expands to: *"The following prompt that we will provide MUST BE executed in background in parallel with all main work streams using the subagents-driven development approach! All work done MUST PRODUCE rock solid evidence covered with hard physical proof(s) that all done is working as expected and as specified without any false results and without any bluff!"* (composes §11.4.20 / §11.4.70 subagent-driven, §11.4.58 / §11.4.103 parallel streams, §11.4.89 background execution, §11.4.5 / §11.4.69 / §11.4.107 captured physical evidence, §11.4 anti-bluff). A second registered action **`REMINDER`** re-surfaces previously-scheduled, CRITICAL, status-UNCERTAIN work: FIRST verify the ACTUAL current status from captured evidence (never assume done or not-done — §11.4.6 no-guessing), THEN act on the delta — report the captured proof if genuinely complete, resume from the exact point if partial, action it NOW if not started, or surface the block per §11.4.66/ §11.4.101 — always producing a status verdict (composes §11.4.6 / §11.4.87 / §11.4.94 / §11.4.97 / §11.4.103 / §11.4.108 / §11.4.130 / §11.4.147).

Severity / handling-priority markers (§11.4.140 extension, 2026-07-16). Three further registered actions tag the REMAINDER of the prompt with a HANDLING PRIORITY (as `BACKGROUND` tags it "run in background"): **`CRITICAL`** (highest-priority / potentially release-blocking — maximum urgency + rigor ahead of lower-priority work, track it, auto-activate §11.4.102 systematic-debugging when an issue is involved, do NOT defer it), **`IMPORTANT`** (high-priority, above routine work but below CRITICAL — track it, full anti-bluff rigor), and **`NOTE`** (capture the remainder as durable CONTEXT — request-history §11.4.208 / §11.4.210 + persistent memory when a non-obvious project fact — applied when relevant, NOT an urgent action unless it contains an explicit action; §11.4.6 record only what the note says, never invent). All six grammar forms work (`CRITICAL: x` ≡ `CRITICAL :: x` ≡ `CRITICAL ---> x` ≡ `/CRITICAL x` ≡ `/DEFAULT::CRITICAL x`). Registering `NOTE` promotes `NOTE:` from an ordinary-prose NO-OP to a registered action (ordinary-prose single-colon examples are now `TODO:` / `WARNING:` / `FIXME:`).

> The system is UNIVERSAL (every CLI agent reads this block via its context carrier per §11.4.35), extensible (new action = new registry row), decoupled + reusable (§11.4.28), and loads out-of-the-box. Classification: universal (§11.4.17). **Canonical authority:** constitution submodule `Constitution.md` §11.4.140. Non-compliance is a release blocker. No escape hatch — no `--skip-action-prefix`, `--ignore-prefix`, `--no-registry`, `--invent-expansion-OK`, `--single-layer-only` flag.
>
> 1 1 4 1 4 1

Token-efficiency mandate (§ . .)

§11.4.141 — Token-efficiency mandate (research-derived + operator mandate, 2026-06-09). Every project worked on by AI coding agents MUST cut token spend (input AND output) toward **30–40% of current (a 60–70% reduction)** WITHOUT degrading quality/performance/safety or breaking any existing mechanism, via a composable, safety-ranked measure set: (1) **prompt-cache the static governance prefix** — the always-loaded governance forms a byte-stable cache-breakpointed prefix with no volatile bytes ahead of it; cache reads cost ~0.1× base input (the dominant cost driver — measured ~170K tokens of governance re-sent every turn, externally corroborated by Claude Code issue #24147); caching is transparent so it removes no rule, weakens no gate, changes no verdict — only billing (PRIMARY, biggest + safest lever); the operative discipline is to **batch governance edits and keep CLAUDE.md + the constitution byte-stable mid-session** — every governance edit busts the prompt cache (BASELINE.md finding); (2) **subagent model-tiering + output-to-file** — mechanical non-judgment work (search/grep/status/doc-export/read-only probes) to a Haiku-class model, the strong model RESERVED for all reasoning/verdicts/fix-design (§11.4.102)/code-review (§11.4.125)/demotion (§11.4.7), large output persisted to a file not an inline 350–520K-token transcript; the cheap model never emits a PASS so §11.4.50 + anti-bluff are untouched; (3) **thin always-loaded INDEX + on-demand detail** — concise index (one line per fix/anchor, EACH carrying the literal `11.4.N` token so propagation gates pass) with the canonical full text kept gate-scanned in `constitution/Constitution.md` and reachable in one hop — a de-duplication realising §11.4.35, never a deletion; (4) **CodeGraph/retrieval-first over full-file loading** (§11.4.78/§11.4.79/§11.4.80); (5) **output-token reduction** — terse status

+ `effort:"low"` on the mechanical allowlist only; (6) **tool-call batching + no re-reads**; (7) **compaction/context-editing for long sessions**. **Mandatory measured proof:** a token-accounting harness measures tokens-per-development-cycle BEFORE vs AFTER on a frozen deterministic workload from the authoritative `usage` object (input/cache_read/cache_creation/output split; NEVER `tiktoken`, NEVER the client-side cost estimate), reproduced N times (§11.4.50), pass = AFTER ≤ 40% of BEFORE OR the measured best-safe reduction with a cited cold-cache reason; the AFTER run MUST show ZERO regression on the pre-build sweep + meta-test mutation sweep + propagation gates + a strong-model reasoning probe + a cache-warm proof (`cache_read_input_tokens > 0`) — cost reduction with quality regression is a §11.4 FAIL. The headline number is the *measured* reduction, never the design estimate (§11.4.6/§11.4.123). No measure may break/degrade any existing mechanism, and the rule is structured so none can. Composes §11.4.5/.6/.20/.40/.50/.58/.69/.70/.78/.79/.80/.103/.106/.123/.125/.128/§12.6/§1.1. Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-141-PROPAGATION` (literal `11.4.141`) + recommended gate `CM-TOKEN-EFFICIENCY` + paired §1.1 mutation (inject a pre-breakpoint volatile token → cache collapses → measured reduction falls below bar → gate FAILs). **Canonical authority:** constitution submodule `Constitution.md` §11.4.141. Non-compliance is a release blocker. No escape hatch — no `--skip-token-efficiency` , `--no-cache-governance` , `--assert-reduction-without-measuring` , `--tier-down-reasoning` , `--inline-all-governance` , `--tiktoken-estimate-OK` flag.

§11.4.142 — Universal code-review mandate — every change reviewed, always, no exception (User mandate, 2026-06-09). Verbatim operator mandate: "ALL changes we do MUST pass through the code review step!!! ALWAYS!!!" EVERY change made to ANY repository this Constitution governs — without exception — MUST pass through an INDEPENDENT code-review step BEFORE it is accepted, committed, or built. NO change class is exempt: source code, fixes, tests, gates, meta-test mutations, documentation, doc-tooling, build/CI scripts, configuration, governance files (Constitution / CLAUDE / AGENTS / QWEN), conductor main-stream edits, sub-agent output, refactors, single-line edits — if a diff exists, it gets an independent review. This is the ABSOLUTE form of §11.4.125 (code-review agent gate after a batch, before pre-build sweep + main build): §11.4.142 strips every implicit scoping §11.4.125's "after all fixes/changes/ implementations are done" phrasing could leave open — no "just a doc edit", no "just a one-liner", no "the author already self-reviewed (§11.4.92)", no "trivial change" carve-out. **Independence is load-bearing** — the reviewer MUST be structurally separated from the author (a dedicated code-review agent, subagent-driven by default per §11.4.70 / §11.4.20, or a distinct human), NEVER the author re-reading their own work; §11.4.92's multi-pass self-evaluation is the AUTHOR-side discipline and PRECEDES (never satisfies) §11.4.142. **The review is itself anti-bluff** (§11.4 / §11.4.1) — a rubber-stamp "looks good" is not a review; it MUST genuinely analyse correctness, safety (no host §12 / data §9 / target-hardware §11.4.133 regression), will-it-really-work (no solve-A-create-B), end-user behaviour (§11.4 / §107), and test-genuineness (§1.1), its conclusions captured evidence per §11.4.5 / §11.4.69, and **it iterates to a clean GO per §11.4.134** — any finding (BLOCKING / nit / warning) re-arms the loop, acceptance only on ZERO new findings + ZERO warnings with rock-solid physical evidence. Honest boundary (§11.4.6): a passing review does NOT replace §11.4.108 four-layer runtime-signature verification nor the §11.4.40 full-suite retest — it is one of MULTIPLE STRONG LAYERS, and the FIRST one every change crosses. Composes §11.4.1 / §11.4.4 / §11.4.5 / §11.4.6 / §11.4.20 / §11.4.40 / §11.4.69 / §11.4.70 / §11.4.92 / §11.4.108 / §11.4.110 / §11.4.125 / §11.4.134 / §107 / §1.1. Classification: universal (§11.4.17) — the consuming project supplies its reviewer-dispatch mechanism + change-acceptance seam (commit wrapper / merge queue / PR gate) per §11.4.35. Propagation gate `CM-COVENANT-114-142-PROPAGATION` (literal `11.4.142`) + recommended gate `CM-EVERY-CHANGE-REVIEWED` (every accepted change carries a fresh independent-review-completed marker for its diff, produced by a reviewer distinct from the author, before acceptance/commit/build) + paired §1.1 mutation (strip the literal → propagation gate FAILs; accept a diff with no independent-review marker → `CM-EVERY-CHANGE-REVIEWED` FAILs). **Canonical authority:** constitution submodule `Constitution.md` §11.4.142. Non-compliance is a release blocker. No escape hatch — no `--skip-review`, `--no-review`, `--trivial-change-exempt`, `--doc-edit-exempt`, `--self-review-suffices`, `--review-after-commit` flag.

§11.4.143 — Real-user-journey mandate for video-streaming-app full-automation tests (User mandate, 2026-06-10). Verbatim operator mandate: "All video streaming apps ... require to choose some title and to press proper UI button to start or resume playing! Proper UI interaction to play exact show with proper content and subtitles is MANDATORY! Without it we just eventually play on 2nd display sample, and that's it mostly! THIS MUST BE ADDED as MANDATORY

RULE regarding testing of any video streaming app in general with full automation tests! Universal in root constitution + a consuming project's extensions." Any full-automation test that asserts a video player / streaming application plays content MUST drive the REAL end-user journey through the app's OWN UI — launch → BROWSE the actual catalog → choose a SPECIFIC title → press the real Play/Resume button → confirm THAT chosen content is genuinely playing, with its correct subtitles, on the intended routing target. A test that bypasses the journey with a sample / demo / built-in-loop clip, a deep-link / `am start -a VIEW` / intent shortcut, a synthetic or pre-staged stream, or any path that does NOT exercise the app's own browse-select-play UI is a §11.4 PASS-bluff at the user-journey layer: it validates ROUTING (that *something* reaches the display) while leaving the user-visible behaviour the operator cares about — "the show I picked actually plays" — unproven (the operator's "we just eventually play on 2nd display sample, and that's it mostly" gap). The mandate (ALL hold): (1) **Real journey, not a shortcut** — launch → catalog browse → specific-title selection → real Play/Resume press through the app's own UI (§11.4.48 UI-driven), NEVER a deep-link / intent / sample / loop-clip shortcut; (2) **Chosen-content confirmation** — the PASS proves the SPECIFIC selected title is playing (not merely that pixels move on the target) via the §11.4.107 liveness battery on the §11.4.136 real-content path + the §11.4.137 subtitle content-correctness oracle for the chosen title's captions; (3) **Non-introspectable UIs use the pixel oracle** — when the app's accessibility hierarchy is blank/unreliable (TV-Compose / leanback / canvas / GL), DRIVE input + ASSERT content via the §11.4.117 CV/OCR pixel oracle, never a hierarchy-only tool; (4) **Login via the credential single-source** (§11.4.10), never hardcoded, never logged; (5) **Honest SKIP, never a faked PASS** — where the autonomous journey is genuinely infeasible (hard human-only login / CAPTCHA, geo-block per §11.4.3, secure-surface pixel-blanking per §11.4.112) the test is `operator_attended` SKIP-with-reason per §11.4.52 + §11.4.3 with a tracked migration item — NEVER a metadata-only / sample-played / routing-only PASS. Honest boundary (§11.4.6): "the routing fired so the title is playing" is a guess — only the chosen-content liveness + subtitle oracle on the real journey proves it. Composes §11.4.48 / §11.4.107 / §11.4.117 / §11.4.136 / §11.4.137 / §11.4.52 / §11.4.3 / §107 / §1.1. Classification: universal (§11.4.17) — the consuming project supplies its concrete app roster, login-credential source, routing target, and UI-driving / pixel-oracle harness per §11.4.35. Propagation gate `CM-COVENANT-114-143-PROPAGATION` (literal `11.4.143`) + recommended gate `CM-VIDEO-REAL-JOURNEY-TEST` (every video-streaming-app playback test drives the real browse-select-play journey + confirms the chosen content + subtitles, or SKIPs-with-reason) + paired §1.1 mutation (replace a real-journey test's browse-select-play path with a deep-link / sample shortcut → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.143. Non-compliance is a release blocker. No escape hatch — no `--sample-playback-ok`, `--skip-real-journey`, `--deep-link-suffices`, `--routing-only-pass`, `--no-title-selection` flag.

§11.4.144 — Tracked/recorded-device availability-following mandate (User mandate, 2026-06-10). Direct operator mandate: every device the project is tracking / following / recording (test / debug / manual-testing device, across every reachable transport — USB / wireless ADB / SSH / serial / network introspection API) MUST be availability-FOLLOWED — its connection state continuously monitored, any drop handled, never silently abandoned and never presented as a continuous recording. The §11.4.128 always-on recorder KNOWS a tracked device is absent (its per-device loop guards on the reachable state) but, lacking a following discipline, merely spins idle — no data captured, no offline event logged, no resume, no escalation — so the recording corpus presents a continuous timeline with a silent hole, a §11.4 PASS-bluff at the recording-integrity layer (it claims continuous capture while no data exists for the gap). On a tracked device leaving its reachable state the system MUST, automatically: (a) **DETECT** the drop and **log an honest offline event** into the recording corpus (§11.4.6 — a silent gap presented as continuous capture is a fabricated-continuity bluff; §11.4.128 — a silent recording gap defeats always-on recording); (b) **WAIT** for the device to return using the project's ALREADY-DEFINED reconnection timings — never invented numbers (§11.4.6), the SAME grace / reconnect / poll budgets the project's recovery path already uses; (c) **RE-ATTACH** — resume recording / tracking the moment the device returns and log an honest online / resume event; (d) **ESCALATE** to the project's sanctioned device-recovery path (per §11.4.69 feature class `device_recovery`) if the device does not return within the defined timeout — through the sanctioned, authorization-gated recovery entry point ONLY, never bypassing its gate and never performing a destructive recovery (e.g. a power-cycle) autonomously without that authorization (§11.4.21 / §11.4.101 — high-blast-radius recovery is gated; while blocked the system keeps following the device, logging the blocked-escalation honestly). A tracked-device drop that produces a silent corpus hole, a never-resumed recording, or a never-escalated permanent absence is the bluff this anchor forbids. Composes §11.4.128 (always-on recording — §11.4.144 closes its drop-handling gap) / §11.4.69 (`device_recovery` sink-side positive evidence) / §11.4.6 (honest offline / online events, reused-not-invented timings) / §11.4.14 (watchdog children reaped on stop) / §11.4.21 + §11.4.101 (gated, non-autonomous escalation). Classification: universal (§11.4.17) — the consuming project supplies its concrete tracking transport, device set, already-defined reconnection timings, and sanctioned recovery entry point per §11.4.35. Propagation gate `CM-COVENANT-114-144-PROPAGATION` (literal `11.4.144`) + recommended gate `CM-DEVICE-AVAILABILITY-FOLLOWED` + paired §1.1 meta-test mutation (strip the watchdog wiring / let an absence go unlogged → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.144. Non-compliance is a release blocker. No escape hatch — no `--skip-availability-following` , `--silent-recording-gap-OK` , `--no-reconnect-wait` , `--invent-reconnect-timing` , `--skip-recovery-escalation` , `--autonomous-power-cycle-OK` flag.

§11.4.145 — Independent multi-angle impact-research per change (User mandate, 2026-06-10). For EVERY fix / change / new feature, INDEPENDENT impact-research agents (subagent-driven §11.4.70/§11.4.20, structurally separate from the author — §11.4.92 self-eval PRECEDES, never satisfies — adversarial "refute-the-change" stance to defeat the documented LLM confirmation-bias failure mode) MUST research the change AND its connected/dependent features across a CLOSED SET OF EIGHT ANGLES — (1) correctness/logic, (2) regression (what existing working feature could break — via call-graph impact enumerating every direct + transitive caller and contract-dependent feature, each mapped to its test), (3) latent/dangerous-code — the recents class (runtime-only failure, interface/ABI/contract mismatch, concurrency/data-race, resource lifecycle), (4) security (taint/injection/secret/unsafe-API), (5) performance (latency/memory/thermal under load vs baseline), (6) host/data/target-hardware safety per §12/§9/§11.4.133, (7) cross-feature interaction (shared state/timing/hardware contention), (8) business-logic conformance (matches the spec AND the connected features' contracts, not merely compiles). Forensic anchor (FACT, project-internal): the recents / 3-button-nav path shipped a latent AIDL-interface-contract mismatch that PASSED code review yet failed at RUNTIME ONLY (ANR on a recents-reaping gesture) because no review angle interrogated the interface contract / concurrency-lifecycle / silently-broken connected feature — the §11.4.108 SOURCE→ARTIFACT→RUNTIME→USER-

VISIBLE gap caught only after the fact; §11.4.145 shifts that discovery LEFT. Each angle MUST name the TOOL(S) used + obtain the required DEPTH (the diff, the full changed unit, its declared contract, the spec section, every connected feature — NEVER the diff lines alone) + emit a captured-evidence conclusion per §11.4.5/ §11.4.69 ("looks fine" forbidden, §11.4.1); a genuinely-N/A angle is recorded `NOT_APPLICABLE: <reason>` per §11.4.6, never silently skipped. Output = a per-change impact-research REPORT (one section per angle: tool + evidence path + verdict) + a single GO/NO-GO that BLOCKS acceptance/commit/build on ANY unmitigated risk in ANY angle; the change is fixed/mitigated + affected angles re-researched, iterating to a CLEAN GO (zero unmitigated risk + zero warning) per §11.4.134 before the §11.4.125 review gate. Honest boundary (§11.4.6): a GO proves cross-angle internal-consistency + bounded blast-radius — it does NOT replace §11.4.108 runtime-signature verification nor §11.4.40 full-suite retest; it is one of MULTIPLE STRONG LAYERS, the FIRST research pass every change crosses. Composes/strengthens §11.4.92/.125/.108/.110/.134/.78/.79/.107/.85/.133/ §12/§9/.99/.6/§1.1. Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-145-PROPAGATION` (literal `11.4.145`) + recommended gate `CM-IMPACT-RESEARCH-PER-CHANGE` + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; accept a change with a report missing an angle or an unmitigated NO-GO → `CM-IMPACT-RESEARCH-PER-CHANGE` FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.145. Non-compliance is a release blocker. No escape hatch — no `--skip-impact-research`, `--single-angle-suffices`, `--author-researches-own-change`, `--diff-skim-OK`, `--no-go-overridable`, `--trivial-change-no-research` flag.

§11.4.146 — Reproduce-first test + same-test-confirms-fix + mandatory extend-to-all-cases workflow (User mandate, 2026-06-10). Every reported problem MUST be handled by a NAMED three-step test workflow — §11.4.146 does NOT re-author its component disciplines, it NAMES + BINDS them into the operator's exact per-defect sequence and adds ONLY the two emphases the components leave implicit. **(STEP 1 — REPRODUCE-FIRST + INVESTIGATE)** before any fix, author the §11.4.43 RED test as a §11.4.115 RED-baseline-on-the-broken-artifact (reproduce the defect on the CURRENT pre-fix artifact, capture defect-present physical evidence per §11.4.5/§11.4.69/§11.4.107); NEW EMPHASIS (D1) that same RED test is ALSO a deliberate investigation instrument per §11.4.102(A) — it MUST gather ADDITIONAL forensic data characterising the defect (triggers, boundaries, input/topology scope, adjacent failure modes) that feeds BOTH the fix design AND the STEP-3 extend scope; a RED test used only as a binary present/absent gate with no characterisation satisfies §11.4.115 but NOT §11.4.146 STEP 1. **(STEP 2 — SAME-TEST-CONFIRMS-FIX)** after the fix the SAME test source (the §11.4.115 polarity switch flipped `RED_MODE=1→0`) confirms the defect ABSENT — RED-on-broken then GREEN-on-fixed, both captured, on a clean target per §11.4.108/§11.4.139, validated-first-after-redeploy per §11.4.130, deterministic per §11.4.50; no separate happy-path test substitutes (the §11.4.43/ §11.4.115 PASS-bluff). **(STEP 3 — EXTEND-TO-ALL-CASES, mandatory per-fix)** NEW EMPHASIS (D2) immediately after STEP 2 the test MUST be FANNED OUT across the full case-space of the SAME functionality — all flows, valid + invalid + boundary (§11.4.85 empty/max/off-by-one) + concurrent (§11.4.85 contention) + failure-injection (§11.4.85 chaos) + topology variants (§11.4.3) — confirming no issue of any kind, with PROVABLE enumerated coverage per §11.4.118 (a listed case-set with per-case outcome, never "no other issues found"); a REQUIRED step, NOT deferred to a release-cycle discovery sweep and NOT reduced to a single guard; each user-visible case carries rock-solid physical evidence per §11.4.123 + is registered into the §11.4.135 standing regression-guard suite; a newly-discovered fan-out defect triggers §11.4.4 test-interrupt + re-enters STEP 1. (ANTI-BLUFF, all steps) every PASS is rock-solid CAPTURED physical evidence per §11.4.123 (§11.4.5/§11.4.69/§11.4.107) — metadata-only / config-only / absence-of-error / grep-without-runtime PASS forbidden (§11.4/§11.4.1); unclear validation method ⇒ deep-research-before-declaring-untestable per §11.4.123/ §11.4.8/§11.4.99; an operator-found defect the green suite missed triggers §11.4.138. Honest boundary (§11.4.6): STEP 3 reduces the unknown-unknown surface but does not prove zero remaining defects (§11.4.118) — un-exercised cases stated as honest gaps, never silently implied clean. Composes §11.4.43/.115/.102/.130/.108/.139/.50/.85/.118/.135/.3/.123/.5/.69/.107/.138/.4/§107/ §1.1. Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-146-PROPAGATION` (literal `11.4.146`) + recommended gate `CM-REPRODUCE-FIRST-THEN-EXTEND` (every closed defect's fix carries a §11.4.115 reproduce-first polarity test with captured defect-characterisation [STEP 1] + its `RED_MODE=0` GREEN confirmation [STEP 2] + an enumerated per-functionality extend case-set registered into the §11.4.135 suite [STEP 3]; a closure with the reproduce→confirm pair but NO enumerated extend case-set FAILs) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; close a defect with only the reproduce→confirm pair and no extend-case-set →

`CM-REPRODUCE-FIRST-THEN-EXTEND` FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.146. Non-compliance is a release blocker. No escape hatch — no `--skip-reproduce-first`, `--fix-without-red`, `--skip-extend-to-all-cases`, `--reproduce-confirm-suffices`, `--defer-extend-to-release-sweep`, `--single-guard-suffices`, `--no-defect-characterisation` flag.

§11.4.147 — Crashed-agent respawn-until-complete + no-work-loss registry mandate (User mandate, 2026-06-10). Verbatim operator intent: "any agent that crashed because of something MUST BE respawned and finish its work at some point! We MUST NOT lose any work, forget about or have it corrupted!" Forensic case study (FACT): dispatched subagents died on TRANSIENT causes (`API Error: Server is temporarily limiting requests` rate-limit killed 5 at once, `API Error: socket connection closed unexpectedly` killed 2, one left PARTIAL edits — two source files written, the dependent SystemUI sliders + doc + test NOT done), each re-dispatched by pure conductor vigilance with NO mechanical guarantee — the moment conductor context is lost / compacted / the conductor itself crashes, an in-flight-but-dead work unit is silently dropped (lost / forgotten / corrupted). Every dispatched agent/subagent MUST be tracked through its full lifecycle so a crash NEVER loses, forgets, or corrupts its work; a crashed agent is NOT a completed agent (§11.4.6 — "it probably finished before dying" is a forbidden guess; abnormal termination is positive evidence the unit is OPEN). The mandate (ALL hold): **(a) durable agent REGISTRY** — every dispatched agent has an append-only machine-readable registry entry (stable id, task + §11.4.58 file-scope, declared output path(s) + §11.4.5/§11.4.69 evidence dir, dispatch timestamp, status from the CLOSED SET `{dispatched | in-flight | crashed | respawned | complete}`), reusing the §11.4.116 sync substrate (JSONL event stream + atomic status snapshot; "an entry with no dispatch-event cannot show `complete` "); the registry is the SINGLE SOURCE OF TRUTH for "is this work owed?" and an unregistered dispatch is itself a violation; **(b) mechanical CRASH-DETECTION + RESPAWN-until-complete** — any abnormal termination (transient rate-limit/socket-close, any non-completion exit, or a terminal error after the runtime's own retries are exhausted) flips the entry to `crashed` + keeps the unit OPEN, and the unit is respawned (fresh agent claims the same id + scope + preserved state) and re-respawned until an agent reaches `complete` ; respawn is the safe/reversible/bounded decision taken autonomously per §11.4.101 (never blocks the loop for a human), transient causes use the project's ALREADY-DEFINED backoff budgets (never invented, §11.4.6), a deterministically-reproducible non-transient terminal error is investigated per §11.4.102 before further respawn (never a blind retry loop); **(c) PARTIAL-STATE preserve → §11.4.84-check → resume-or-clean-restart** — the crashed agent's uncommitted edits + output doc are PRESERVED (never silently discarded → lost, never blindly committed → corruption), the respawn runs the §11.4.84 quiescence check on the preserved tree (every modified file accounted-for vs scope, no mutation/ `// always pass` / `_mutated_*` residue, no half-written/torn artifact), then EITHER RESUMES idempotently when it passes OR CLEAN-RESTARTS from a known-good base (§9.2 pre-op backup, reversible §11.4.101) when inconsistent — nothing lost, nothing corrupted; per-agent `git worktree` isolation (§11.4.58 L4 / §11.4.84) keeps the partial tree from contaminating other streams; **(d) "crash ≠ done" COMPLETION criterion** — a unit is DONE only when an agent reaches `complete` with its required captured evidence/output landed (§11.4.5/§11.4.69 + the §11.4.116 verdict-carries-evidence-path rule); the endless-loop done-condition (§11.4.87/§11.4.94/§11.4.97/§11.4.126) MUST NOT read satisfied while any entry is `dispatched` / `in-flight` / `crashed` / `respawned` -not-yet- `complete` , the zero-idle survey (§11.4.94) treats every non- `complete` entry as an OPEN item to reclaim, and a registry showing `complete` without landed evidence is a §11.4 PASS-bluff at the agent-lifecycle layer. Honest boundary (§11.4.6): the registry + respawn guarantee work is not lost/corrupted, NOT that it is correct — the respawned output still crosses §11.4.108 + §11.4.125/§11.4.142 review + §11.4.40 retest; §11.4.147 is the durability layer beneath those, the agent-side analogue of §11.4.144 (same detect→wait/backoff→re-attach/respawn→escalate shape). Composes §11.4.6/.58/.84/.87/.94/.97/.101/.116/.102/.108/.125/.142/.40/.128/.144/ §9.2/§1.1. Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-147-PROPAGATION` (literal `11.4.147`) + recommended gate `CM-CRASHED-AGENT-RESPAWN-TRACKED` + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; mark a `crashed` entry as the loop's done-condition `complete` without landed evidence, OR drop a dispatched agent without a registry entry → `CM-CRASHED-AGENT-RESPAWN-TRACKED` FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.147. Non-compliance is a release blocker. No escape hatch — no `--skip-agent-registry`, `--crash-equals-done`, `--no-respawn`, `--discard-partial-state`, `--blind-commit-partial`, `--forget-dead-agent`, `--loop-done-with-crashed-entries` flag.

§11.4.148 — Workable-item integrity (status+type+id) + comprehensive structured description + bidirectional external-tracker sync + BLOCKED unblock-choices mandate (User mandate, 2026-06-10). BINDS + STRENGTHENS the workable-item discipline (§11.4.15 status / §11.4.16 type / §11.4.54 id / §11.4.21 Operator-blocked / §11.4.91 description-clarity / §11.4.93 + §11.4.95 SQLite SSoT / §11.4.106 docs_chain) into ONE integrity contract spanning DB ↔ docs ↔ external tracker, adding the operator's three emphases: **(D1) no item without a valid status + valid type + stable id — on ALL three surfaces** (an item missing any of the three in the DB, the rendered docs, OR the external tracker FAILs the validator — release-blocker; the id is the cross-surface binding key); **(D2) comprehensive structured description per item** — WHAT it is (§11.4.91 ≥6-word/≥40-char clear meaning) + HOW it manifests + HOW to reproduce (§11.4.115/.146) + ACCEPTANCE CRITERIA (§11.4.123/.69 captured-evidence verdict), in the §11.4.93 DB `description` column AND the docs AND the tracker, never a stub/§11.4.91-anti-pattern fragment; **(D3) BLOCKED items carry WHY + enumerated unblock CHOICES** — tightens §11.4.21: every `Operator-blocked` item (incl. `Blocked` / `BLOCKED` documented alias normalised to canonical `Operator-blocked`, never a silent fork §11.4.6) MUST enumerate the closed list of decisions/actions that would unblock it (`[A]…·[B]…·[C]…`, mirroring §11.4.66's 2–4-option shape); a BLOCKED item with no enumerated choices FAILs; **(D4) regular never-missed bidirectional DB↔docs↔tracker sync** — the §11.4.93/.95 git-tracked SQLite DB is the SSoT, docs + tracker DERIVED; full sync (`db-to-md` / `md-to-db` byte-identical round-trip §11.4.93 + tracker push) runs regularly + on every change, §11.4.86 drift-proof fingerprint (sha256 of the sorted item keyset, NOT mtime) gates freshness, §11.4.106 docs_chain-bound so drift is mechanically caught not vigilance-dependent; **(D5) generic external-tracker push** carries statuses (collapsed onto the tracker's native set per §11.4.33/.112 when fixed, precise value preserved in a header, never lost), types, assignee (§11.4.104 participant handle, UNSET defaults from a project env var, never hardcoded/logged §11.4.10), and sub-tasks, IDEMPOTENT (dry-run-then-real, match-by-stable-key `[<ID>]` prefix / id custom-field, present⇒UPDATE/absent⇒CREATE, rate-limited, credential-redacted, sink-side `created=N updated=M failed=0` proof §11.4.69), the MACHINERY project-agnostic §11.4.28

(consumer registers its tracker/list/field-map at runtime). Anti-bluff §11.4: every sync + validator pass carries captured evidence §11.4.5/.69; honest boundary §11.4.6 — guarantees well-formed items + agreeing surfaces, NOT that the underlying work is correct (still crosses §11.4.108/.40/.123). Composes §11.4.15/.16/.21/.33/.34/.54/.66/.86/.91/.93/.95/.104/.106/.112/.123/.10/.28/.69/.6/ §1.1. Classification: universal (§11.4.17) — consumer supplies its DB path, id prefix, tracker service + list/board id + field map + default-assignee env var, docs_chain context per §11.4.35. Propagation gate `CM-COVENANT-114-148-PROPAGATION` (literal `11.4.148`) + recommended gates `CM-ITEM-INTEGRITY-STATUS-TYPE-ID` / `CM-ITEM-COMPREHENSIVE-DESCRIPTION` / `CM-BLOCKED-UNBLOCK-CHOICES` / `CM-TRACKER-SYNC-IDEMPOTENT` + paired §1.1 mutation (gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.148. Non-compliance is a release blocker. No escape hatch — no `--item-without-status`, `--item-without-type`, `--item-without-id`, `--stub-description-OK`, `--blocked-without-choices`, `--skip-tracker-sync`, `--one-way-sync-OK`, `--non-idempotent-tracker-push`, `--hardcode-assignee` flag.

§11.4.149 — Per-workable-item testing-diary mandate (User mandate, 2026-06-10). Every workable item MUST carry an append-only TESTING DIARY — a chronological record of every test run against it — distinct from §11.4.93 `item_history` (lifecycle STATE transitions) + §11.4.55 Reopens (reopen cycles): the diary records TEST EXECUTIONS (which may or may not change status). The mandate (ALL hold): **(a)** a `test_diary` table additive to the §11.4.93/.95 SQLite SSoT, one row per run soft-keyed to the item id, capturing ISO-8601-UTC `date_time` + `tested_by` (closed set `User|Operator|AI-agent|HelixQA`) + `result` (§11.4.45 `PASS|FAIL|SKIP` + detail) + in-depth `observations` (long markdown, facts or explicit `UNCONFIRMED:` / `PENDING_FORENSICS:` §11.4.6, the background a future fix needs) + `action_taken` + `status_changed` +from/to (did this run change status + WHY) + §11.4.69 `evidence_path` + `feature_class`, with a SCHEMA CONSTRAINT making a `PASS` row WITHOUT a non-empty evidence path impossible (a PASS-bluff rejected by the schema itself); **(b)** BOTH an in-depth diary doc + a derived at-a-glance summary VIEW (total/pass/fail/skip + last-verdict + last-run + status-changes + distinct testers/feature-classes, DERIVED never duplicated §11.4.93) feeding §11.4.132 risk-ordering; **(c)** four-format per-item `Diary` / `Diary_Summary` exports §11.4.65 + §11.4.86-drift-proof-fingerprinted (sha256 of the sorted diary keyset) + §11.4.106 docs_chain-bound so a diary row not exported/pushed FAILs a gate (never-missed); **(d)** external-tracker SUB-TASK model — the item task's description stays the item (§11.4.148 D2, NOT polluted with diary text), each run a CHILD sub-task (type `task`) with a `{TODO|In-progress|Completed}` lifecycle (collapsed per §11.4.33/.112), observations + action + an Evidence: line (path not raw artefact §11.4.10/.13) + a diary-entry idempotency key, reusing the §11.4.148 D5 idempotent rate-limited credential-redacted sink-side-proven push, mapper project-agnostic §11.4.28; **(e)** MINIMAL-LLM deterministic bash/Go tooling (zero LLM in the data path — the `observations` prose is authored by whoever ran the test; tooling only stores/renders/pushes/validates), 100% test-covered §11.4.27 (unit + integration-against-the-real-tracker with an honest §11.4.3 SKIP-when-token-absent never a faked PASS + export + HelixQA Challenge + paired §1.1 mutation) so a diary PASS-bluff is mechanically impossible. Honest boundary §11.4.6 — the diary guarantees a complete auditable per-item test-history, NOT that any single run's verdict is correct (rests on its own §11.4.69/.107/.123 evidence). Composes §11.4.6/.27/.45/.50/.55/.65/.69/.86/.93/.95/.106/.107/.123/.132/.148/.10/.13/.28/.33/.112/ §1.1. Classification: universal (§11.4.17) — consumer supplies its DB path, diary doc layout, export formats, tracker sub-task field map, docs_chain context per §11.4.35. Propagation gate `CM-COVENANT-114-149-PROPAGATION` (literal `11.4.149`) + recommended gates `CM-TEST-DIARY-SYNC` / `CM-DIARY-PASS-REQUIRES-EVIDENCE` + paired §1.1 mutation (gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.149. Non-compliance is a release blocker. No escape hatch — no `--skip-testing-diary`, `--diary-without-evidence`, `--no-diary-summary`, `--diary-without-tracker-subtask`, `--llm-in-diary-data-path`, `--diary-export-optional` flag.

§11.4.150 — Mandatory deep multi-angle web research per change/issue, before declaring fixed or structural (User mandate, 2026-06-11). Verbatim operator mandate: "For every single issue we fix or improvement we make — besides tight systematic-debugging, fixing, code review by independent agents, comprehensive tests — we MUST ALWAYS do deep web research from various angles! No matter how big/small/simple/complex, dig deep the internet for articles, technical documentation, APIs, open-source code — EVERYTHING that can help make the best possible solution OR confirm we don't have a more serious problem we're unaware of! We MUST do everything possible to STOP the constant issue-reopening and finally start closing items as fixed+working! Do this ALWAYS in parallel with the main work stream!" For EVERY fix / improvement / change / closure — no matter how big, small, simple, or complex — the agent MUST, IN ADDITION to tight systematic-debugging (§11.4.102), the fix, independent-agent code review (§11.4.125 / §11.4.142), and comprehensive multi-layer tests (§11.4.4(b) / §11.4.40), perform a DOCUMENTED **deep multi-angle web research pass** digging the internet from VARIOUS angles — articles, official + vendor technical documentation, API references, standards, issue trackers, maintainer guidance, reusable open-source code — to BOTH (i) discover the best possible solution AND (ii) confirm there is NOT a more serious underlying problem the team is unaware of. Forensic case study (FACT, 2026-06-11): a subtitle-on-secondary-display goal verdicted `Won't-fix: structurally-impossible` (§11.4.112 secure-surface pixel-blanking) was OVERTURNED by a deep multi-angle research pass that surfaced the real OCR-of-the-PRIMARY-display path — a "we can't" became a shipped capability; the canonical anti-pattern of a structural verdict reached for want of research. The mandate (ALL hold): **(A) No closure-as-fixed/ structural WITHOUT a documented deep-research pass** — no item may be marked `Fixed` / `Implemented` / `Completed` (§11.4.33) OR classified `structurally-impossible` won't-fix (§11.4.112) until a deep multi-angle pass is documented; §11.4.150 makes §11.4.8's pass UNCONDITIONAL (every item, however trivial, at the closure AND structural-verdict gates). **(B) Multiple angles, not a single lookup** — ≥ 2 genuinely-distinct angles (official-docs / standards / known-bug-trackers / alternative-approach / failure-mode / security / performance / platform-constraint / open-source-precedent) so a single-source confirmation-bias miss (§11.4.145) cannot pass; a one-link drive-by is NOT deep multi-angle. **(C) Confirm-no-bigger-problem, not just find-a-fix** — explicitly seek evidence the fix does not mask a deeper defect AND the closure/verdict is genuinely safe; "we found nothing worse" requires the enumerated search (§11.4.118). **(D) Latest-source + cited** — LATEST authoritative versions per §11.4.99 (never training-data/ memory), each cited by URL + access date in the research artefact AND the closure commit footer (`Deep-research <date>: <urls>` OR the literal `NO external solution found — original work` per §11.4.8). **(E) Reopen-breaking is the PURPOSE** — STOP the constant Fixed↔Reopened churn (§11.4.34 / §11.4.55); a reopen whose root cause a deep pass would have surfaced is a §11.4.150 miss; composes §11.4.7 (demotion-evidence) + §11.4.112 (structural verdict now REQUIRES the cited-authorities pass). **(F) ALWAYS in parallel with the main stream** — background subagent-driven (§11.4.70 / §11.4.20 / §11.4.103 / §11.4.89) concurrent with main fix/build/test work, NEVER serialising it or stalling the loop (§11.4.94 / §11.4.97 / §11.4.101 / §11.4.126). **(G) Apply ASAP to every workable item** — retroactively + going forward across the §11.4.93 / §11.4.95 SSoT; items closed without a documented pass are re-audited in the §11.4.40 / §11.4.42 release-gate sweep. Honest boundary (§11.4.6): the pass reduces the unknown-unknown surface + breaks the most common reopen causes — it does NOT prove zero remaining defects (§11.4.118) and does NOT replace §11.4.108 runtime-signature verification, §11.4.125 / §11.4.142 review, or §11.4.40 retest; it is the research layer every fix/closure/structural-verdict additionally crosses, and "we probably don't have a bigger problem" without the enumerated multi-angle search is a guess (§11.4.6), never a finding. Composes §11.4.8 / §11.4.99 / §11.4.123 / §11.4.118 / §11.4.145 / §11.4.125 / §11.4.142 / §11.4.7 / §11.4.112 / §11.4.34 / §11.4.55 / §11.4.70 / §11.4.20 / §11.4.89 / §11.4.103 / §11.4.40 / §11.4.42 / §11.4.93 / §11.4.95 / §11.4.6 / §1.1. Classification: universal (§11.4.17) — the consuming project supplies its concrete research corpora, angle set, item tracker, and closure-commit-footer convention per §11.4.35. Propagation gate `CM-COVENANT-114-150-PROPAGATION` (literal `11.4.150`) + recommended gate `CM-DEEP-RESEARCH-PER-ISSUE` (every closed/structural-verdicted item carries a documented multi-angle deep-research artefact + cited-source closure footer, run in parallel, before the closure/structural verdict is accepted) + paired §1.1 mutation (strip the literal → propagation gate FAILs; close an item or reach a `structural` verdict with no documented multi-angle research artefact / cited footer → `CM-DEEP-RESEARCH-PER-ISSUE` FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.150. Non-compliance is a release blocker. No escape hatch — no `--skip-deep-research`, `--single-source-suffices`, `--trivial-change-no-research`, `--close-without-research`, `--structural-without-research`, `--serialise-research`, `--research-from-memory-OK` flag.

§11.4.151 — Project-prefixed release-tag/version-naming mandate (User mandate, 2026-06-12). Verbatim operator mandate: "Every release tag and version name we create — on the main repository and on every Submodule we own — MUST be prefixed with the project's release prefix, e.g. `myproject-1.0.0-dev-0.0.1`. The prefix MUST come from `HELIX_RELEASE_PREFIX` in our `.env` if it is set, otherwise from the lowercased project root directory name. The SAME prefix MUST be used across the main repo and all owned Submodules in one release so a release is greppable across every repository." Every release tag AND every version name created on the main repository AND on every owned-by-us submodule (§11.4.28) MUST be prefixed with the project's release prefix, form `<PREFIX>-<version>` (canonical example: `myproject-1.0.0-dev-0.0.1`), so a release is identifiable + greppable across every repository it spans (one `git tag -l '<PREFIX>-*'` enumerates the whole release surface). **Prefix resolution order (closed-set, deterministic — §11.4.6 no-guessing):** (1) `HELIX_RELEASE_PREFIX` from the project's `.env` — authoritative when set; `.env` is git-ignored per §11.4.30 and the variable is documented in the tracked `.env.example` (a §11.4.77 re-obtain mechanism, never committed); (2) fallback = the lowercased snake_case form of the project root directory name (no spaces) per §11.4.29 — used whenever the env var is unset/empty, so a prefix is ALWAYS resolvable from the checkout with zero operator input. **The prefix MUST be IDENTICAL across the main repo and all owned submodules within a single release** — a release that tags the main repo `<PREFIX>-1.2.0` while tagging an owned submodule with a different/unprefixed value is a §11.4.151 violation (the cross-repo grep no longer enumerates the release). Version codes increment monotonically within the prefix (`<PREFIX>-…-0.0.1` → `<PREFIX>-…-0.0.2` → …), never reset, never skipped. Honest boundary (§11.4.6): the prefix guarantees a release is identifiable + uniform across every repository, NOT that its contents are correct — the tag is still created only after the §11.4.40 full-suite retest GREEN and reaches every upstream via the §11.4.113 merge-onto-latest-main path (NEVER a force-push), fanned out per §2.1. Composes §2.1 / §11.4.29 / §11.4.30 / §11.4.40 / §11.4.113 / §11.4.126 (the release-scope terminal condition is a published, prefixed tag). Classification: universal (§11.4.17) — the consuming project supplies its concrete prefix value + the `HELIX_RELEASE_PREFIX` env var per §11.4.35. Propagation gate `CM-COVENANT-114-151-PROPAGATION` (literal `11.4.151`) + recommended gate `CM-RELEASE-PREFIX-NAMING` (every release tag/version on the main repo + every owned submodule carries the resolved `<PREFIX>-` prefix, identical across the release) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; create an unprefixed or differing-prefix release tag → `CM-RELEASE-PREFIX-NAMING` FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.151. Non-compliance is a release blocker. No escape hatch — no `--no-release-prefix`, `--unprefixed-tag`, `--prefix-optional`, `--differing-submodule-prefix` flag.

§11.4.152 — Crashlytics-recorded-data continuous monitoring + systematic-debug + regression-test-coverage mandate (User mandate, 2026-06-13). Verbatim operator intent: "For every project that has Firebase Crashlytics enabled /

wired, we MUST continuously monitor ALL of the Crashlytics-recorded data — crashes, ANRs, performance traces, and non-fatals — systematically debug each, fix and improve, and cover everything with validation and verification tests. This MUST be checked regularly, with no false results and no bluff of any kind!" Every project that has Firebase Crashlytics enabled/wired (SDK linked into a shipping artifact, crash+non-fatal+ANR reporting active) MUST treat the Crashlytics console as a first-class captured-evidence channel from real end-user devices and continuously process every datum it records — an open Crashlytics issue with a green test suite is a §11.4 PASS-bluff at the field-telemetry layer (a real user hitting a broken feature while the suite reports green). STRENGTHENS+COMPLETES §11.4.47: §11.4.47 owns the periodic REVIEW + dedup + Issue-creation pass that SURFACES Crashlytics/Analytics/Performance findings; §11.4.152 owns what happens to each surfaced item AFTER — the systematic-debug → fix/improve → regression-test-coverage lifecycle that drives it to a proven regression-immune closure (§11.4.47 finds it; §11.4.152 fixes it + proves it stays fixed; both mandatory, neither substitutes). The mandate (ALL hold): (1) **continuous monitoring of ALL four surfaces** — fatal crashes, ANRs, performance traces/regressions, AND non-fatals (skipping any one is a PASS-bluff; non-fatals are the silent class — every catch/fallback/recovered-error path is a real degraded UX and MUST be triaged, not ignored because the app did not crash); (2) **systematic-debugging of each (reproduce-before-fix)** per §11.4.102 Iron Law + §11.4.115 RED-baseline-on-the-broken-artifact (the recorded stacktrace/ANR-trace/non-fatal-context IS the §11.4.5/.69 captured defect-present evidence) BEFORE the fix — a fix without a prior reproducing test is a §11.4.43/.123 violation; unreproducible ⇒ deep-research-before-declaring-untestable §11.4.123/.150; (3) **a fix/improvement for every confirmed issue** against its proven root cause (§11.4.9/.43/.108), "acknowledged in console" is not a fix, muting a recurring issue without a tracked rationale is a §11.4.6/.90 violation; (4) **validation+verification regression-test coverage per closed issue** — in the SAME commit as the fix register a permanent §11.4.135 regression guard (§11.4.115 polarity test: `RED_MODE=1` captures the recorded defect on the pre-fix artifact, `RED_MODE=0` standing GREEN guard asserting ABSENT) exercising the same user-reachable path the stacktrace identifies, with rock-solid captured evidence §11.4.5/.69/.107/.123 + a closure log citing the console issue id/URL + root-cause + fix commit + the validation/verification test paths; a Crashlytics issue marked resolved WITHOUT a falsifiable regression test is FORBIDDEN (the silent-recurrence vector, §11.4.138 operator-escape class applied to field telemetry); (5) **regular cadence** reusing the §11.4.47 five-trigger set (pre-build/pre-flash/pre-distribute/pre-tag blocking; daily + post-deployment-burn-in non-blocking) — a one-time sweep never repeated is a violation; (6) **no false results, no bluff** (§11.4/.1/.6 — "no new issues" requires the enumerated monitored-surfaces+window §11.4.118; "fixed" requires the RED→GREEN flip §11.4.115/.130; a guard whose §1.1 mutation does not FAIL it is a bluff gate). Honest boundary §11.4.6: processing every recorded issue breaks the silent-recurrence vector — it does NOT prove zero remaining field defects nor replace §11.4.108/.40; an issue is "closed" only when its guard is GREEN on a clean artifact, never when the console mark is flipped. Classification: universal (§11.4.17) — the consuming project supplies its Crashlytics handle, console-access credential (§11.4.10, never logged), severity table, and per-issue closure-log path per §11.4.35; the reference project-level instantiation is a consuming project's §6.O (Crashlytics-Resolved Issue Coverage — per-issue validation test + Challenge test + `.lava-ci-evidence/crashlytics-resolved/<date>-<slug>.md` closure log) + §6.AC (Comprehensive Non-Fatal Telemetry — every catch/fallback records a non-fatal with triage context). Composes §11.4/.1/.5/.6/.9/.34/.40/.43/.47/.69/.90/.102/.107/.108/.115/.118/.123/.130/.135/.138/.150/ §1.1. Propagation gate `CM-COVENANT-114-152-PROPAGATION` (literal `11.4.152`) + recommended gate `CM-CRASHLYTICS-ISSUE-FULLY-COVERED` (every closed Crashlytics issue carries a systematic-debug root-cause record + a registered §11.4.135 regression guard + a closure-log entry citing the console issue id/URL + the validation/verification test paths; a console-resolved issue with no falsifiable regression guard FAILs) + paired §1.1 meta-test mutation (gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.152. Non-compliance is a release blocker. No escape hatch — no `--skip-crashlytics-monitoring`, `--monitor-crashes-only`, `--skip-non-fatals`, `--resolve-without-regression-test`, `--console-mark-is-fixed`, `--monitor-once`, `--mute-without-rationale` flag.

§11.4.153 — Comprehensive per-feature Status + Status_Summary document set with mandatory video-recording confirmation (User mandate, 2026-06-15). Every project MUST maintain, under `docs/features/`, a comprehensive feature Status document set (`Status.md` + its §11.4.56 `Status_Summary.md` companion) enumerating EVERY system component, EVERY client app/binary/ surface (TUI / CLI / Web / desktop / mobile / API / gRPC / library / submodule / infrastructure), and EVERY feature — including features ported from any incorporated CLI-agent / submodule catalogue (§11.4.74) — with NO single feature left out. (1) **Total categorized coverage** — per-feature table organized Component→Category, reconciled against the actual codebase (a code-present feature missing from the table, or a row with no code, is a §11.4.6/§11.4.118 bluff). (2) **Per-feature fields** — Component / Feature / Category / Implementation / Wiring (genuinely end-user-reachable, not merely compiled, §11.4.108) / Real-use / Tests-coverage (four-layer §11.4.4(b)) / Validation (PASS / FAIL / SKIP / PENDING_FORENSICS / OPERATOR-BLOCKED per §11.4.45) / **Video-recording confirmation** (path to the real-use video or honest gap marker). (3) **Mandatory per-feature real-use video** — every user-visible "confirmed" claim backed by a recorded real-use video of a genuine end-user scenario (real prompts → real LLM/service responses → real results), NEVER a frozen/stale frame (§11.4.107), NEVER faked/mocked/demo-loop/bluff response or LLM error passed off as success (§11.4.2/§11.4.5), stored at the project-declared recording path (§11.4.35); a confirmed row with no real video or a bluff video = §11.4 PASS-bluff; autonomous-infeasible ⇒ honest §11.4.3/§11.4.52 SKIP + tracked migration item, NEVER a faked confirmation. (4) **Video-analysis remediation loop** — every video analysed; any defect it surfaces triggers §11.4.102 systematic-debug → fix → §11.4.146 retest → re-record → clean GO per §11.4.134 before "confirmed" (a video exposing a broken feature is a §11.4.4 test-interrupt). (5) **Always-in-sync** — §11.4.45-class roster/corpus-backed, §11.4.106 docs_chain-bound + §11.4.86 drift-proof fingerprint (sha256 of sorted feature-key roster AND sorted video-artefact roster, NOT mtime), re-syncs out-of-the-box; stale = violation. (6) **Four-format export** — HTML + PDF + **DOCX** (this doc class ADDS DOCX to the §11.4.65 HTML+PDF set; other classes unchanged), in sync per §11.4.60. (7) Follows §11.4.44/.45/.56/.57/.59/.60. (8) **MP4 format REQUIRED.** All video confirmations MUST be in `.mp4` format (H.264). Window-specific capture ONLY (§11.4.159(A)). Vision validation REQUIRED (§11.4.159(D)). `.cast` files are supplementary only. Honest boundary §11.4.6 — guarantees a complete video-confirmed always-synced ledger, NOT per-feature correctness (rests on each row's §11.4.69/.107/.123 evidence) and NOT a §11.4.40 retest substitute. Classification: universal (§11.4.17). Composes §11.4.2/.5/.44/.45/.52/.56/.57/.59/.60/.65/.86/.102/.106/.107/.108/.118/.123/.134/.146/§1.1. Propagation gate `CM-COVENANT-114-153-PROPAGATION` (literal `11.4.153`) + recommended gates `CM-FEATURE-STATUS-COMPLETE` + `CM-FEATURE-STATUS-VIDEO-CONFIRMED` + paired §1.1 mutation.

Canonical authority: constitution submodule `Constitution.md` §11.4.153. Non-compliance is a release blocker. No escape hatch — no `--skip-feature-status`, `--feature-without-video`, `--frozen-video-OK`, `--bluff-response-video-OK`, `--skip-video-analysis`, `--feature-ledger-incomplete-OK`, `--no-docx-export`, `--allow-feature-ledger-drift` flag.

§11.4.154 — Window-scoped capture + fresh-corpus rotation for feature/QA recordings (User mandate, 2026-06-15). Verbatim: "recording you perform does record the window containing apps and services, not the whole desktop or monitor screen!" + "All old recording files MUST BE removed when new one starts!" Refines §11.4.2/.5/.107/.153 recording discipline with two capture-hygiene invariants. **(A) Window-scoped, NOT whole-screen** — every feature/QA video MUST capture ONLY the window/surface of the app/service under test (GUI window / CLI-TUI terminal pane / web tab-viewport / device-emulator-simulator frame), NEVER the whole desktop/monitor or unrelated windows; whole-desktop capture leaks operator-private content (§11.4.10/.83), dilutes the §11.4.107 liveness/freeze oracle, and breaks the §11.4.137 OCR/ROI content oracle. Target the window/region by stable identity (window id/title, device serial, browser context, tmux target) per §11.4.111 — never a fixed full-screen device index. Platform genuinely cannot capture below whole-screen ⇒ honest §11.4.3 SKIP + tracked migration item, never a whole-screen pass-off. **(B) Fresh-corpus rotation** — when a new recording run for a scope begins, the agent's OWN prior in-scope stale recordings at the raw recording path MUST be removed FIRST so the live corpus reflects the current run (§11.4.107 not-stale + §11.4.86 roster-freshness). Honest boundary (§11.4.6 + §9.2): "remove old" = the agent's own prior recordings for the SAME scope/project ONLY — NEVER another project's/scope's/operator-authored files; uncertain ⇒ surface, don't delete (§11.4.122); committed `docs/qa/<run-id>/` evidence (§11.4.83) is the durable record, NOT rotated. Classification: universal (§11.4.17). Composes §11.4.2/.5/.10/.83/.86/.107/.111/.122/.128/.137/.153/§9.2/.6. Propagation gate `CM-COVENANT-114-154-PROPAGATION` (literal `11.4.154`) + recommended gate `CM-WINDOW-SCOPED-FRESH-CORPUS-RECORDING` + paired §1.1 mutation.

Canonical authority: constitution submodule `Constitution.md` §11.4.154. Non-compliance is a release blocker. No escape hatch — no `--whole-screen-capture-OK`, `--skip-window-scope`, `--keep-stale-recordings`, `--no-corpus-rotation`, `--full-desktop-recording` flag. **(C) MP4 auto-conversion**

REQUIRED. Any `.cast` file produced MUST be auto-converted to `.mp4` via `agg` + `ffmpeg` immediately after capture. The `.mp4` is the primary evidence; `.cast` is supplementary only.

§11.4.155 — Project-name-prefixed feature/QA recording filenames (User mandate, 2026-06-15). Verbatim: "All recorded videos MUST START with prefix: the PROJECT NAME (ALWAYS USE THE PROJECT NAME). Project name MUST be obtained according to the constitution's own project-name resolution." Every recorded video the project produces — every feature/QA real-use recording (§11.4.153), every window-scoped capture (§11.4.154), every always-on device recording (§11.4.128), every raw or curated artefact at the project-declared recording path (§11.4.35) + the committed `docs/qa/<run-id>/` trail (§11.4.83) — MUST have a filename that STARTS WITH the PROJECT-NAME prefix, ALWAYS; an unprefixed recording is a §11.4.155 violation (a multi-project/scope corpus on one host per §11.4.128/.103 becomes un-greppable + un-attributable — the §11.4.151 identify-and-grep failure on the recording-corpus axis). **Prefix resolution (closed-set, deterministic — §11.4.6, IDENTICAL to §11.4.151):** (1) `HELIX_RELEASE_PREFIX` from `.env` (authoritative, git-ignored §11.4.30, documented in tracked `.env.example` §11.4.77) else (2) lowercased snake_case project-root dir name §11.4.29 — ALWAYS resolvable, zero operator input. SAME prefix for EVERY recording in a checkout so `ls '<PREFIX>---'*` enumerates the corpus; canonical form `<PREFIX>---<feature-or-scope>---<run-id>.<ext>` (`---` delimits the prefix unambiguously); MUST equal the §11.4.151-resolved release-tag prefix (divergence is itself a §11.4.155 violation — one project, one name). Honest boundary (§11.4.6): the prefix guarantees attribution + greppability, NOT content validity (still §11.4.107/.137/.153) and does NOT relax §11.4.154's window-scope/rotation (rotation removes the agent's OWN `<PREFIX>---*` only; foreign/operator files surfaced not deleted §11.4.122/§9.2). Classification: universal (§11.4.17). Composes §11.4.151/.128/.153/.154/.111/.83/.6/.29/.30/.35/.77/.86/§1.1. Propagation gate `CM-COVENANT-114-155-PROPAGATION` (literal `11.4.155`) + recommended gate `CM-RECORDING-PROJECT-NAME-PREFIX` + paired §1.1 mutation.

Canonical authority: constitution submodule `Constitution.md` §11.4.155. Non-compliance is a release blocker. No escape hatch — no `--no-recording-prefix`, `--recording-without-project-name`, `--unprefixed-recording`, `--prefix-optional-for-recording`, `--differing-recording-prefix` flag.

§11.4.156 — All CI/CD automation (GitHub Actions / GitLab pipelines / equivalents) MUST be disabled (User mandate, 2026-06-15). Verbatim operator mandate: "Any GitHub actions or GitLab pipelines MUST BE disabled! Add this critical mandatory rule / mandatory constraint into the root constitution Submodule, commit and fetch all its changes to all upstreams and make sure we respect and follow this rule (we do apply it) ASAP!!!" Every repository this Constitution governs — main repo, this constitution submodule, every owned + nested submodule we author and push — MUST ship with ALL server-side CI/CD automation DISABLED: no push to any owned upstream may trigger a GitHub Actions run, GitLab pipeline, or equivalent (Jenkins/CircleCI/Travis/Drone/Woodpecker/Bitbucket/Azure, any `on: push` / `schedule` / `workflow_dispatch`). GENERALISES + makes ABSOLUTE the §11.4.75 Layer-5 posture (remote CI DISABLED, workflow preserved at a `…disabled-local-only` non- `.yml` name a provider ignores) across ALL governed repos; enforcement migrates to the LOCAL §11.4.75 git-hook ritual + §11.4.40 pre-tag sweep, never a remote runner. ALL hold: **(A)** zero active `.github/workflows/*.yml|yaml` / `.gitlab-ci.yml` / `.gitlab/**` / equivalent at the ROOT of any governed repo/submodule (the only place a provider executes); **(B)** "disabled" = a push triggers ZERO runs — delete OR rename to a non-trigger name (§11.4.75 `.disabled` / `.disabled-local-only`); a live- `on:` + `if:false` workflow still queues runs, NOT compliant; **(C)** scope = repos we author+push — vendored/third-party nested configs below the root (AOSP `external/**` , `prebuilts/**` , vendored submodules) are INERT (a provider never runs a non-root config), OUT of scope, MUST NOT be mass-edited (§11.4.29 vendor-exempt); test = "does a push to OUR upstream trigger a run?" yes⇒disable, inert⇒document+leave (§11.4.6 verify-not-assume); **(D)** no new CI may be added (release blocker); **(E)** pre-push verify `git ls-files | grep -E '^\.github/workflows/.*\.ya?ml$|^\.gitlab-ci\.yml$'` empty for authored repos, §11.4.109-class PreToolUse guard + gate enforce mechanically. Honest boundary (§11.4.6): file-level disabling stops FILE-triggered runs, NOT provider-side server settings (org-default required workflows, branch-protection required checks, provider scheduled exports) — the operator turns those off; the agent documents what it cannot reach, never claims unachieved completeness. Composes §11.4.75/.29/.6/.40/.42/.109/.113/§2.1. Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-156-PROPAGATION` (literal `11.4.156`) + recommended gate `CM-NO-ACTIVE-CI` + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; add a root `.github/workflows/x.yml` to an authored repo → `CM-NO-ACTIVE-CI` FAILs). **Canonical authority:** constitution submodule `Constitution.md` §11.4.156. Non-compliance is a release blocker. No escape hatch — no `--allow-ci`, `--enable-workflow`, `--keep-pipeline`, `--remote-ci-OK`, `--ci-exempt` flag.

§11.4.157 — GEMINI.md maintained in lockstep with CLAUDE.md / AGENTS.md / QWEN.md (User mandate, 2026-06-15). Verbatim operator mandate: "Make sure with CLAUDE.md, AGENTS.md, QWEN.md we maintain GEMINI.md too! Add this mandatory fact / rule to root constitution Submodule we are inheriting / extending - CONSITUTION.md, CLAUDE.md, QWEN.md, AGENTS.md, GEMINI.md and other related relevant files!" Forensic FACT (2026-06-15): when §11.4.156/§11.4.157 were authored, Constitution/CLAUDE/AGENTS/QWEN carried the family through §11.4.155 but GEMINI.md had silently drifted to §11.4.141 — 14 mandates (§11.4.142–155) never propagated — a §11.4 propagation-bluff (a Gemini-CLI agent reads a stale Constitution). GEMINI.md is a FIRST-CLASS governance context carrier EQUAL to CLAUDE.md/AGENTS.md/QWEN.md, never optional/best-effort. ALL hold: **(A)** five-carrier lockstep — no governance change is complete until GEMINI.md carries it alongside the other three mirrors; GEMINI.md is added to the §11.4.26 propagation + cross-reference set explicitly; **(B)** no silent drift — GEMINI.md lagging the other mirrors' highest rule is a §11.4.157 violation (§11.4.65-class), back-fill required; **(C)** equal status — GEMINI.md restates the SAME literal `11.4.N` anchors the propagation gates require (§11.4.35), fleet count INCLUDES GEMINI.md; **(D)** consumer projects' own CLAUDE/AGENTS/QWEN/GEMINI bind too (§11.4.35). Honest boundary (§11.4.6): the §11.4.142–155 GEMINI.md back-fill is a tracked release-blocking remediation; claiming GEMINI.md "in sync" while the back-fill is incomplete is itself a §11.4.157 violation. Composes §11.4.26/.35/.17/.44/.65/.140/.156/§1.1. Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-157-PROPAGATION` (literal `11.4.157`, GEMINI.md INCLUDED) + recommended gate `CM-GEMINI-MD-LOCKSTEP` + paired §1.1 meta-test mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.157. Non-compliance is a release blocker. No escape hatch — no `--skip-gemini-md`, `--gemini-optional`, `--gemini-lag-OK`, `--four-carrier-suffices` flag.

§11.4.158 — Intensive all-feature/flow/edge-case video-recording + read-the-screen content-verification mandate (User mandate, 2026-06-16). Every project MUST be covered by intensive automated testing that exercises + RECORDS every feature/flow/use-case/edge-case (valid/invalid/boundary/concurrent/failure-injection §11.4.85), each recording showing the feature GENUINELY WORKING with REAL results (real prompts→real responses→real outputs §11.4.153) and NO false/simulated/stale/frozen result (§11.4.2/.5/.107) — a recording showing an error/bluff/non-working feature is a FINDING (→§11.4.153(4)/§11.4.4 fix→retest→re-record), never a confirmation. The testing System MUST ACTUALLY READ every shown log line / message / UI label / dialog / toast / status text and VERIFY it is a genuine working result (OCR/ROI §11.4.117/.137 + confidence floor for pixel surfaces; direct capture for terminal/log; §11.4.107(10) self-validated golden-good/golden-bad analyzer) — "a video was produced" is NOT evidence, "the System read the screen + confirmed a genuine result" is. HelixQA (§11.4.27) MUST drive this exercise→record→read→score pass, PASS only on a read-confirmed genuine result + captured artefact path (§11.4.69). Default recording save path = `$HOME/Downloads` (host user's home Downloads, resolved at runtime never hardcoded) unless a project declares an override per §11.4.35; §11.4.155 project-prefix + §11.4.154 window-scope/rotation + §11.4.128 git-ignored raw corpus + §11.4.83 curated docs/qa evidence apply. **Vision analysis MANDATORY for every recording** — after every recording, the agent MUST read the terminal output / video content and verify the feature actually works. A recording without a vision-confirmed verdict is a §11.4.158 violation. Composes §11.4.2/.5/.25/.27/.52/.69/.83/.85/.107/.108/.117/.118/.128/.137/.138/.153/.154/.155/§1.1. Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-158-PROPAGATION` (literal `11.4.158`) + recommended gates `CM-INTENSIVE-RECORDING-COVERAGE` + `CM-RECORDING-CONTENT-READ-VERIFIED` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.158. Non-compliance is a release blocker. No escape hatch — no `--skip-recording-coverage`, `--video-without-content-read`, `--happy-path-recording-suffices`, `--recording-path-anywhere`, `--unread-recording-OK`, `--skip-helixqa-read` flag.

§11.4.159 — Mandatory window-specific video recording + vision validation mandate (User mandate, 2026-06-20). Every feature test, validation, verification, challenge, and QA session that produces video evidence MUST comply with ALL of the following: **(A) Window-specific recording ONLY** — every video MUST record ONLY the target application window (Terminal pane, TUI app, browser tab, emulator frame), NEVER the whole desktop/monitor screen; use macOS `screencapture -l<window_id>`, Linux `xdotool` + `ffmpeg`, or equivalent; whole-desktop capture leaks operator-private content (§11.4.10), dilutes the §11.4.107 liveness oracle, and breaks §11.4.137 OCR/ROI; platform genuinely cannot capture below whole-screen => honest §11.4.3 SKIP + tracked migration item. **(B) MP4 format REQUIRED** — `.mp4` (H.264, `movflags +faststart`, `pix_fmt yuv420p`); `.cast` files are supplementary only; auto-convert via `agg` + `ffmpeg` per §11.4.154(C). **(C) Project-name prefix REQUIRED** — filename MUST start with project name in snake_case from `HELIX_RELEASE_PREFIX` in `.env` (§11.4.151) or lowercased project root dir name (§11.4.29); format `<project_name>-<feature>-<timestamp>.mp4`; unprefixed = violation. **(D) Mandatory vision validation** — after EVERY recording, the agent MUST read the terminal output / video content and verify: (i) feature ACTUALLY WORKS, (ii) LLM responses are REAL, (iii) all tests show PASS, (iv) no "TODO implement" / "simulate" / "for now" patterns, (v) output demonstrates end-user working feature; MUST produce a verdict (PASS/FAIL) with evidence path (§11.4.69). **(E) Terminal window cleanup** — after each recording, dismiss/close ONLY the Terminal window used for that recording (window-specific close via `osascript` / `xdotool`); MUST NOT close windows belonging to other processes. **(F) Real results ONLY** — every recording MUST show REAL working features; errors/empty output/simulated responses trigger fix→retest→re-record. **(G) Re-runnable evidence** — the command shown MUST be re-runnable to produce the same results. **(H) Fresh-corpus rotation** — remove agent's own prior in-scope stale recordings FIRST (§11.4.154). **(I) Content verification MANDATORY — not duration-based** — the value of a recording is NOT its duration but its CONTENT; a recording MUST demonstrate the ACTUAL feature being used with REAL results; before accepting ANY recording, the agent MUST verify: expected output patterns ARE present (e.g., test PASS lines, API response data, feature-specific output), the feature ACTUALLY WORKS as demonstrated (not just "something ran"), LLM responses are REAL content (not simulated, not placeholder, not empty), every claim of "working" is backed by visible evidence; a 5-second recording showing a feature working correctly is MORE valuable than a 60-second recording of empty terminal; duration is NOT a proxy for quality. **(J) Expected-content specification REQUIRED** — before recording, the agent MUST specify what content SHOULD appear (expected patterns, expected test results, expected API responses); after recording, the agent MUST verify these patterns ARE present; if expected content is MISSING, the recording is REJECTED regardless of duration. **(K) Content-verification recording workflow** — mandated workflow: (1) SPECIFY expected content patterns, (2) RECORD the feature execution, (3) EXTRACT all text from the recording, (4) VERIFY expected patterns are present, (5) CHECK for simulated/placeholder content, (6) ACCEPT only if ALL patterns found AND zero bluffs detected, (7) REJECT and re-record if ANY pattern missing or bluff detected. **(L) Root cause analysis REQUIRED for rejected recordings** — when a recording is rejected (missing expected content, bluff detected, or empty capture), the agent MUST investigate WHY before re-recording per §11.4.102; determine the root cause (timing issue, wrong command, tool failure, etc.) and fix it; simply re-recording without understanding WHY the first attempt failed is a §11.4.159 violation. **(M) Real-time monitoring RECOMMENDED** — for complex features, use real-time monitoring that analyzes output DURING recording (not after); this catches issues immediately and allows corrective action before the recording completes. Classification: universal (§11.4.17). Composes §11.4.2/.3/.5/.10/.29/.69/.83/.107/.111/.128/.137/.151/.153/.154/.155/.158/§1.1. Propagation gate `CM-COVENANT-114-159-PROPAGATION` (literal `11.4.159`) + recommended gate `CM-WINDOW-VIDEO-VALIDATED` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.159. Non-compliance is a release blocker. No escape hatch — no `--whole-screen-ok`, `--cast-only`, `--skip-vision-validation`, `--no-cleanup`, `--simulated-recording-ok`, `--unprefixed-recording` flag.

§11.4.160 — Vision-verified recording + HelixQA bridge mandate (User mandate, 2026-06-21). Compact summary: every video recording for feature/QA evidence MUST be processed through a vision/OCR pipeline that reads on-screen content and confirms expected results BEFORE acceptance; the recording system MUST provide a bridge feeding captured frames to HelixQA's test infrastructure (or equivalent) for automated read-the-screen verification against SPECIFY-phase expected patterns (§11.4.159(J)); capture frames at ≤5s intervals, run OCR/vision with self-validated analyzer (§11.4.107(10)), compare text against patterns, produce per-frame PASS/FAIL with evidence path, surface failures immediately for re-record (§11.4.159(L)). Project-configured parameters per §11.4.35. Honest boundary: does NOT replace §11.4.5 quality analysis nor §11.4.108 runtime-signature; FLAG_SECURE uses §11.4.117 proxy oracle. Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-160-PROPAGATION` + recommended gate `CM-VISION-VERIFIED-RECORDING-BRIDGE` + paired §1.1 mutation.

§11.4.161 — Rootless container runtime mandate (User mandate, 2026-06-21). Compact summary: every project MUST use Podman rootless (or equivalent) for ALL containerized workloads — rootful Docker/sudo FORBIDDEN unless §11.4.112 platform constraint documented; `vasic-digital/containers` Submodule

(§11.4.76) sole orchestration layer via `pkg/boot` / `pkg/compose` / `pkg/health`; extend upstream per §11.4.74 for missing capabilities; integration tests boot infra on-demand. Honest boundary: does NOT replace §11.4.10 credentials nor §12.3 hygiene. Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-161-PROPAGATION` + recommended gate `CM-ROOTLESS-CONTAINER-RUNTIME` + paired §1.1 mutation.

§11.4.162 — OpenDesign UI design system mandate (User mandate, 2026-06-21). Compact summary: every project with user-facing interfaces MUST use OpenDesign (https://github.com/nexu-io/open-design) as mandatory UI design system — NOT ad-hoc CSS; install as dependency, use token system for all color (light+dark from brand assets), typography, spacing, component tokens; extend upstream (§11.4.74) for unsupported patterns; every UI component ships light+dark + token-diff regression tests via visual regression; no element overlap/font collision/label overlay. Honest boundary: does NOT replace §11.4.27 functional testing nor §11.4.48/.49 UI testing. Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-162-PROPAGATION` + recommended gate `CM-OPENDESIGN-UI-SYSTEM` + paired §1.1 mutation.

§11.4.163 — Universal Media Validation (User mandate). Every recorded artifact MUST pass MEDIA VALIDATION pipeline — OCR/transcription/text parsing, compare vs SPECIFY-phase patterns (§11.4.159(J)), self-validated analyzer (§11.4.107(10)), structured verdict, pinpoint on FAIL, REAL-TIME trigger option. Paired §1.1 mutation on golden-bad fixture. Classification: universal. Propagation gate `CM-COVENANT-114-163-PROPAGATION` + paired §1.1 mutation.

§11.4.164 — Universal Auto-Propagation Hook (User mandate). Every constitutional pull triggers `post_update_hook.sh` — detects changed files, registers skills/MCP/hooks/scripts, logs summary. Consumer invokes after every pull. Classification: universal. Propagation gate `CM-COVENANT-114-164-PROPAGATION` + paired §1.1 mutation.

§11.4.165 — Universal Independent Verification Agent (User mandate). Every code/media/docs/config output passes independent verifier (§11.4.70/.20), zero-finding GO (§11.4.134). CODE: review+build+mutations+runtime-signature. MEDIA: pipeline+genuine+format. DOCS: exports+revision. CONFIG: syntax+schema+leaks. Self-validated §1.1 mutation. Classification: universal. Propagation gate `CM-COVENANT-114-165-PROPAGATION` + paired §1.1 mutation.

§11.4.166 — REPEALED (operator decision, 2026-06-22). The Universal Semgrep static-analysis mandate is repealed — Semgrep is NO LONGER mandatory. Static analysis remains encouraged, not mandated. Scaffolding, `submodules/semgrep`, MCP/pre-commit/PATH wiring, docs_chain `semgrep_status` context, and the `CM-COVENANT-114-166-PROPAGATION` / `CM-SEMGREP-WIRED` gates removed. Anchor 11.4.166 retired. See constitution `Constitution.md` §11.4.166.

§11.4.167 — Big-work-item feature work-stream lifecycle (User mandate, 2026-06-23). Every BIG work item (new feature OR drastic/large fix) MUST be its own isolated feature work-stream — own sibling project copy, own `feature/<slug>` branch, own per-feature flashable builds + feature tags, SEPARATE from trunk until operator manually approves after testing, trunk merged INTO every stream regularly. (A) one big item = one sibling project copy (`<project>_<slug>`); small changes stay on trunk. (B) disk-feasible CoW/reflink clone MANDATORY (`cp -a --reflink=auto` on btrfs/XFS/ZFS/APFS, else `git worktree`); naive deep copies FORBIDDEN; per-stream `out/` + shared ccache; ~0-disk verified (§11.4.69) before relied on (§11.4.6). (C) own branch + greppable tags (§11.4.151 `<base>-feat-<slug>`), NEVER merged to trunk until §11.4.40 GREEN on clean target → §11.4.41 merge-first → §11.4.113 ff-only. (D) trunk merged INTO every live stream per trunk-tag/daily (merge, never rebase). (E) feature branch + tag cascades to EVERY touched submodule (§11.4.113). (F) single-builder build QUEUE (§11.4.58/.12.7/.12.8) + finite-device queue (§11.4.119); idle `out/` GC'd (§11.4.77). (G) every change crosses full gauntlet (§11.4.142/.125/.134/.145 + anti-bluff). (H) stream registry (§11.4.116/.147) + atomic `.partial` → rename + §11.4.84 resume + §9.2 retire-backup; bounded §12/.12.6. Honest boundary (§11.4.6): isolation+disk-feasibility+greppable-identity+merge-discipline, NOT correctness (still §11.4.108/.40). Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-167-PROPAGATION` (literal `11.4.167`) + gates `CM-FEATURE-WORKSTREAM-COW-CLONE` / `-NO-MERGE-UNTIL-APPROVED` / `-TRUNK-SYNC-CADENCE` / `-SUBMODULE-CASCADE` + paired §1.1 mutations. **Canonical authority:** constitution `Constitution.md` §11.4.167.

§11.4.168 — Exported-document independent content + textual + full-visual validation mandate (User mandate, 2026-06-23). Every generated/exported document (HTML/PDF/DOCX/any §11.4.65 export-scope format) MUST pass INDEPENDENT validation by a reviewer structurally separate from the generator (§11.4.70/.142, NEVER author self-check), on BOTH source AND every exported artifact, ALWAYS, across THREE layers: (1) CONTENT — export faithfully carries source intent+data, nothing dropped/truncated/garbled; (2) TEXTUAL — human-readable, NO raw markup/diagram-source (Mermaid `gantt` / `graph` / `flowchart` / `sequenceDiagram` etc.)/unrendered code-fence leaking as body text (the exact 2026-06-23 raw-gantt-in-PDF failure); (3) FULL VISUAL — diagrams render as IMAGES not source, layout intact, no overlap/cut-off/garble, verified by RENDERING the export (`pdftotext` catches raw source + `pdfimages` confirms rendered images + `pdftoppm` → OCR confirms human-readable content) with captured evidence §11.4.5/.69/.107. Anti-bluff: rubber-stamp ≠ validation; analyzer self-validated golden-good/golden-bad §11.4.107(10) (golden-bad seeded with raw `gantt` MUST FAIL); iterate to clean GO §11.4.134. Forensic FACT (2026-06-23): a "green" Mermaid fix shipped raw gantt source as readable text into user PDFs because validation grepped HTML for `class="mermaid"` and NEVER checked the exported PDF — operator caught it (§11.4.138). Honest boundary (§11.4.6): confirms the exported artifact is faithful/readable/visually-rendered — does NOT prove source content is correct (§11.4.142/.135) nor replace the §11.4.65 mtime freshness gate (this is the READABILITY/FIDELITY layer the mtime gate cannot see); FLAG_SECURE/un-rasterisable surfaces document the gap §11.4.112 + §11.4.117 proxy, never a faked PASS. Classification: universal (§11.4.17). Composes §11.4.65/.73/.107/.117/.135/.138/.142/.159/.163/.165/§1.1. Propagation gate `CM-COVENANT-114-168-PROPAGATION` (literal `11.4.168`) + recommended gate `CM-EXPORTED-DOC-VISUALLY-VALIDATED` + paired §1.1 mutation. **Canonical authority:** constitution `Constitution.md` §11.4.168.

§11.4.169 — Mandatory comprehensive test-type coverage with anti-bluff captured evidence (User mandate, 2026-06-25). Every project MUST cover a CLOSED enumerated test-type set, each to as-close-to-100% as the domain permits, every PASS citing rock-solid captured PHYSICAL evidence (§11.4.5/.69/.107), zero false results/zero bluff (§11.4/.1/.6): (1) unit (only layer mocks allowed §11.4.27(A)); (2) integration (real System, infra via containers submodule §11.4.76); (3) e2e; (4) full-automation (autonomous, re-runnable, no manual step §11.4.25/.52/.98, deterministic §11.4.50); (5) Challenges (challenges submodule banks §11.4.27(B)); (6) HelixQA (helix_qa submodule — written test banks/suites per app/service/platform + comprehensive fully-autonomous QA sessions §11.4.27); (7) DDoS/load-flood (clean refuse or graceful degrade); (8) security (authz/injection/secret-leak §11.4.10/crypto/CVE/default-deny + security submodule); (9) stress+chaos (load+contention+boundary+failure-injection, categorised recovery §11.4.85); (10) concurrency/atomicity (no lost updates, transactional atomicity, idempotent replay); (11) race-condition/deadlock (race detector + lock-order analysis; no blocking op in shared-lock region §1); (12) memory (leak census + peak-RSS ceiling + 24h/N-iter soak, no unbounded growth); (13) benchmarking/performance (p50/p95/p99 + throughput vs recorded baseline). The ONLY permitted absence is an honest §11.4.3 SKIP-with-reason, never a silent gap; four-layer §11.4.4(b) enforcement + coverage ledger (§11.4.25/.52); missing required type or OPERATOR_ATTENDED_ONLY = release blocker. Strict enumerated expansion of §11.4.27. Composes §11.4.4/.5/.6/.25/.27/.50/.52/.69/.76/.85/.98/.107/.123/.135/.146/§1.1. Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-169-PROPAGATION` (literal `11.4.169`) + recommended gate `CM-MANDATORY-TEST-TYPES-COVERED` + paired §1.1 mutation. **Canonical authority:** constitution `Constitution.md` §11.4.169.

§11.4.170 — Device-independent host-side rendered-UI visual-proof mandate (User mandate, 2026-06-25). Compact summary: every change to ANY user-facing UI surface (screen/component/view/widget/layout/theme on any platform — Android-Compose, iOS-SwiftUI/UIKit, web, desktop, TUI) MUST be proven by DEVICE-INDEPENDENT host-side RENDERED PIXELS before it may be claimed correct — the real component rendered to a PNG ON THE HOST (NO device, NO emulator, NO running app) via a host-render harness (Compose→Roborazzi/Paparazzi on the JVM; web→Playwright/Storybook snapshot; SwiftUI→snapshot-testing; equivalent per stack), for EVERY relevant screen×state×{light,dark} theme, with the PNG validated by BOTH (i) golden image-diff (per-pixel/perceptual PASS/FAIL) AND (ii) an OCR/vision oracle reading rendered text+labels+control bounds to assert legibility + NO overlap / NO label-over-label / NO clipping / NO off-screen control / NO collapsed-or-giant-unbounded widget (the §11.4.162 layout invariants PROVEN on real pixels). Forensic FACT (2026-06-25): UI work was "validated" by VALUE-EQUALITY unit tests (a hex colour string / an sp/dp size integer / a design-token field == a constant) that NEVER RENDERED A PIXEL — GREEN while the operator opened a broken unbounded giant-button screen. A value/token-equality / property-assertion UI test is FORBIDDEN as the PROOF a UI is correct (it may SUPPLEMENT, NEVER SUBSTITUTE the rendered-pixel proof); a "green" UI suite with no rendered-pixel artefact for the changed surface is a §11.4 PASS-bluff at the visual layer. Self-validated golden-good/golden-bad analyzer §11.4.107(10), thresholds calibrated on the project's own fixtures §11.4.6. "device offline" is

NEVER a valid skip — host-render IS the device-independent path; honest §11.4.3 SKIP only where the platform genuinely has no host-render harness (tracked migration item, never a value-only fallback). Honest boundary §11.4.6: host-render proves the COMPONENT renders correctly device-independently per commit — it does NOT prove the running app on real hardware (that remains §11.4.153/.158/.159/.160's live-device recording + read-the-screen layer, which §11.4.170 COMPLEMENTS not replaces — both mandatory), does NOT replace §11.4.162 OpenDesign token/theme governance (§11.4.170 PROVES the rendered RESULT of those tokens), DISTINCT from §11.4.168 exported-document visual validation. Classification: universal (§11.4.17). Composes §11.4.3/.5/.27/.40/.69/.107/.108/.117/.153/.158/.159/.160/.162/.163/.168/.169/§1.1. Propagation gate `CM-COVENANT-114-170-PROPAGATION` (literal `11.4.170`) + recommended gate `CM-HOST-RENDERED-UI-VISUAL-PROOF` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.170. Non-compliance is a release blocker. No escape hatch — no `--value-assertion-proves-ui`, `--skip-rendered-pixel-proof`, `--token-equality-suffices`, `--device-offline-skip-visual`, `--single-theme-only`, `--unvalidated-image-diff-OK`, `--no-ocr-layout-oracle` flag.

When in doubt

- Read `constitution/Constitution.md` for the canonical text of every rule.
- Cross-reference the sibling carriers (`CLAUDE.md` / `AGENTS.md` / `QWEN.md`) — they carry the same universal content for their respective agents.
- If still unclear, ask the operator. Do NOT guess. Do NOT bluff (§11.4.6).

§11.4.171 — Mandatory comprehensive human-readable workable-item descriptions (User mandate, 2026-06-29). Every workable item (ATM-NNN ticket) in the project's workable-items database, Issues.md, Fixed.md, Issues_Summary.md, Fixed_Summary.md, and ALL derived documents MUST carry a comprehensive human-readable description that explains the item in plain, non-technical language suitable for non-developers (project managers, stakeholders, team leads, business partners). The description MUST be 5-7 sentences minimum, written so that ANY person — regardless of technical background — can understand: (a) WHAT the item is (the feature, fix, or task in simple terms); (b) WHY it matters (the user/business value); (c) HOW it works (the mechanism, without code references); (d) WHO benefits (the end user or stakeholder); (e) WHAT the expected outcome is (the measurable result). The description is MANDATORY — an item without a human-readable description is a §11.4 violation at the documentation layer. Descriptions MUST NOT use jargon, acronyms without expansion, code references, or technical implementation details as the primary explanation. Technical details MAY appear as supplementary context AFTER the plain-language explanation. The `description` column in the SQLite database (`docs/workable_items.db`) is the SINGLE SOURCE OF TRUTH for item descriptions; all derived documents (Issues.md, Fixed.md, summaries, exports) MUST carry the SAME description text. Pre-build gate `CM-HUMAN-READABLE-DESCRIPTION` (when implemented) walks every item in the DB and asserts the `description` field exists, is non-empty, and is ≥ 50 characters. Paired §1.1 meta-test mutation strips a description → gate FAILs. Classification: universal (§11.4.17). Composes §11.4.15 (item-status tracking), §11.4.16 (item-type tracking), §11.4.19 (column alignment), §11.4.91 (summary clarity), §11.4.93 (SQLite SSoT), §11.4.148 (item integrity), §11.4.65 (universal export). Propagation gate `CM-COVENANT-114-171-PROPAGATION` (literal `11.4.171`) + paired §1.1 mutation.

§11.4.172 — Mandatory production-readiness planning with realistic timeline projection (User mandate, 2026-06-29). Every project under this Constitution MUST maintain a living production-readiness planning document that: (a) provides REALISTIC timeline projections for all product milestones (MVP, beta, production-ready) based on actual development velocity data (items completed per week, reopening rate, blocking dependencies); (b) identifies ALL danger zones, weaknesses, and risk areas with mitigation strategies; (c) accounts for HW delivery timelines, licensing delays, and external dependency risks; (d) defines the critical path and identifies which items can slip without affecting MVP; (e) includes workstation/build-capacity analysis and optimization recommendations; (f) is updated at least monthly or whenever the item count changes by ≥ 10%; (g) is cross-referenced against the workable-items database for accuracy. The planning document MUST live at `docs/research/production_planning_<YYYYMMDD>/ANALYSIS.md` with HTML+PDF exports per §11.4.65. Honest boundary (§11.4.6): projections are estimates based on measured velocity — they guarantee the methodology is sound, NOT that specific dates will be met. Classification: universal (§11.4.17). Composes §11.4.6 (no-guessing — projections based on data, not wishful thinking), §11.4.40 (full retest before tag), §11.4.42 (iteration discipline), §11.4.93 (SQLite SSoT), §11.4.108 (four-layer verification). Propagation gate `CM-COVENANT-114-172-PROPAGATION` (literal `11.4.172`) + paired §1.1 mutation.

§11.4.173 — Containerized + distributed build mandate (User mandate, 2026-06-29). EVERY build of EVERY component (source compile, artifact/package/installer/container-image production, codegen, asset render — for ANY language/platform: Go, Android/Gradle, desktop, web, native, firmware) MUST run INSIDE a specialized build container provisioned via the `digital.vasic.containers` submodule (§11.4.76) — NEVER on the bare host. The build containers MUST be DISTRIBUTED to the designated remote build host(s) (e.g. `thinker.local`) via the SAME containers-submodule distribution mechanism the infra uses (§11.4.76 Distributor / remote compose over SSH, §11.4.161 rootless), so the build EXECUTES on the remote build host (offloading the developer/main host); once the build completes the produced artifacts MUST be brought BACK to the originating main host (scp/rsync/volume copy) for use/flashing/distribution. Building outside a container, or on the bare host, is FORBIDDEN — a release blocker (the "works on my machine" / unreproducible-build class §11.4.76 exists to prevent). The build-host target + build-container definitions are config-injected (§11.4.28, never hardcoded in the submodule); a missing build-container capability is added by EXTENDING the containers submodule upstream (§11.4.74), never by an ad-hoc host build. Honest boundary (§11.4.6): the containerized+distributed build guarantees reproducibility + host-isolation + capacity offload — it does NOT replace the §11.4.40 full-suite retest, §11.4.108 four-layer artifact→runtime verification, or §11.4.38 installable-asset evidence (those still run against the brought-back artifact). Classification: universal (§11.4.17). Composes §11.4.76 (containers submodule — sole orchestration layer), §11.4.161 (rootless runtime), §11.4.74 (extend-don't-reimplement), §11.4.28 (config injection / decoupling), §11.4.24 (build-resource stats), §11.4.82 (iteration speedup — persistent caches in the container), §11.4.121 (no-commit-while-build-writes-artifacts), §11.4.38 (installable-asset evidence on the brought-back artifact), §11.4.108 (artifact→runtime verification), §12.6 (host memory ceiling — offloading the build preserves the main host). Propagation gate `CM-COVENANT-114-173-PROPAGATION` (literal `11.4.173`) + recommended gate `CM-CONTAINERIZED-DISTRIBUTED-BUILD` (every build runs via the containers submodule on the remote build host; a bare-host build is detected + FAILs) + paired §1.1 mutation. **Canonical authority:** constitution submodule

`Constitution.md` §11.4.173. Non-compliance is a release blocker. No escape hatch — no `--build-on-host`, `--skip-container-build`, `--local-build-ok`, `--no-distributed-build`, `--bare-host-build` flag.

§11.4.174 — Shared-host process-ownership verification mandate (User mandate, 2026-06-29). On any shared host where other projects/users/agents may run concurrently, every inspection of or action on a process, build, daemon, port, lock, or system-state signal MUST first positively VERIFY the target is OURS before diagnosing or acting. Attribute a process to us ONLY by a strong discriminator — cwd inside this project's tree, argv naming our scripts/artifacts/repo path, a PID we launched + recorded (§11.4.147), or a lock/port path under our tree — NEVER a loose `pgrep java/gradle/node` name match (it matches every project on the host). NEVER kill/signal/restart a process not positively ours (other projects' gradle/java/podman/zig are off-limits — operator-workload-safety, §12/§11.4.133); the §12.8 daemon-kill remediation MUST be scoped to OUR daemons. When our work waits on host contention from another project's process, surface it (§11.4.66/§11.4.101) — block-don't-break. Every pgrep/ps filter MUST exclude self-matches + other-project matches. Forensic anchor (FACT 2026-06-29): a 200%-CPU java+gradlew on the shared host was a separate `catalogizer` project's gradle test (cwd=Projects/catalogizer), nearly mis-attributed to our build. Composes §11.4.6/§11.4.66/§11.4.101/§11.4.147/§12/§12.8/§11.4.133. Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-174-PROPAGATION` (literal `11.4.174`) + recommended gate `CM-PROCESS-OWNERSHIP-VERIFIED` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.174. Non-compliance is a release blocker. No escape hatch — no `--assume-process-ours`, `--loose-pgrep-ok`, `--kill-any-matching-process`, `--skip-ownership-check` flag.

§11.4.176 — Conflict-free multi-track work-division + exactly-once claim registry + capability-aware deadlock-proof device-lock (User mandate, 2026-07-02). Two mandatory multi-track arbitration layers, conflict-free by construction: (A) exactly-once work-item/logical-group claim registry (extends §11.4.147/§11.4.116; disjoint file-scope §11.4.58 L3 + worktree/CoW §11.4.167; no codebase loss §9.2/§11.4.113/§11.4.84/§11.4.41); (B) capability-aware deadlock-proof device-lock (§11.4.119 single-owner; all-or-nothing + non-blocking + TTL-reap breaks all 4 Coffman conditions; append-only JSONL + atomic snapshot §11.4.116); (C) universal/decoupled (§11.4.17/§11.4.28/§11.4.35) + evidence-based resource-tuning (§11.4.24/§11.4.6, never guessed). Classification: universal

(§11.4.17). Gates CM-WORK-DIVISION-EXCLUSIVE-CLAIM + CM-DEVICE-LOCK-DEADLOCK-FREE + CM-COVENANT-114-176-PROPAGATION + paired §1.1 mutation. Composes §11.4.6 §11.4.17 §11.4.24 §11.4.28 §11.4.35 §11.4.41 §11.4.58 §11.4.84 §11.4.85 §11.4.103 §11.4.113 §11.4.116 §11.4.119 §11.4.147 §11.4.167 §12.6 §9.2 §11.4.176.

§12.11 — Maximal dynamic resource utilization for containerized builds (User mandate, 2026-07-03). The containerized build path (its OWN cgroup + own MemoryMax → OOM-kills ITSELF, never user.slice) MUST use the maximal SAFE fraction of host capacity, computed DYNAMICALLY per host: auto-detect nproc/MemTotal/RLIMIT_NPROC (never hardcode, §11.4.6), reserve explicit host headroom (cores+RAM), scale -j against BOTH the memory model AND the process rlimit (privileged nproc-raise for full -j; EAGAIN-safe -j otherwise), never break/bottleneck/wedge. NO fixed low -j for the containerized path. The §12.6 60% ceiling + bare-metal -j cap REMAIN, SCOPED to user.slice-resident work — §12.11 does NOT weaken §12.6, it scopes it. Classification: universal (§11.4.17). Gate CM-MAXRES-DYNAMIC-BUILD + CM-COVENANT-12-11-PROPAGATION + paired §1.1 mutation. Composes §12.1 §12.2 §12.3 §12.6 §12.10 §11.4.6 §11.4.24 §11.4.133 §9.2 §11.4.35 §12.11.

§11.4.177 — Developer-tooling project-decoupling + invocation-directory operation (User mandate, 2026-07-04). Compact summary: no project-specific script / hook / alias / binary may be wired (copied / symlinked / PATH-exported / aliased) into a global or shared developer-tooling PATH; shared tooling MUST be project-agnostic and operate on the INVOCATION directory (cwd / explicit path / config arg), NEVER a hardcoded project path; a project-specific helper lives in + runs from its own repo, a genuinely-reusable helper is promoted to the shared toolkit ONLY after being stripped of every project-specific assumption (paths, device serials, package names, region endpoints) + taking its target from the invocation context; a project script reachable on the global/shared PATH is a decoupling violation of equal severity to importing project-specific code into a shared submodule (§11.4.28); forensic FACT 2026-07-04 — a project cwd-hook symlinked into the global Claude-Toolkit PATH re-coupled every alias to one hardcoded checkout; honest boundary §11.4.6 — decoupling guarantees project-agnostic, not correct (still needs the tool's own tests). Classification: universal (§11.4.17). Composes §11.4.28/.29/.35/.11/.6. Propagation gate `CM-COVENANT-114-177-PROPAGATION` (literal `11.4.177`) + recommended gate `CM-TOOLING-PROJECT-DECOUPLED` + paired §1.1 mutation. **Canonical authority:** constitution submodule

`Constitution.md` §11.4.177. Non-compliance is a release blocker. No escape hatch — no `--global-symlink-ok`, `--hardcode-project-path`, `--wire-into-shared-path` flag.

§11.4.178 — Track-qualified identity for parallel work streams (session-name collision ban) (User mandate, 2026-07-04). Compact summary: parallel work streams sharing hardware / a git backing store / a tmux-session namespace / a lock namespace / any single-owner resource MUST be addressed by a TRACK-QUALIFIED identity (`<project>__<track>__<role>`), NEVER a bare directory basename; multiple checkouts/clones/worktrees sharing basename `atmosphere` collapse into one session/lock/log key + cross-wire tmux targets, lock files, device claims, log dirs (observed collision 2026-07-04); every registry row / session name / lock path / log dir / device claim for a ≥2-concurrent-claimant resource MUST carry the track-qualified key, a bare-basename identity for such a resource is a violation; honest boundary §11.4.6 — a unique identity prevents cross-wiring, it does not itself arbitrate ownership (that is §11.4.119 + **§11.4.176** (the multi-track work-division + exactly-once-claim + device-lock arbitration anchor)). Classification: universal (§11.4.17). Composes §11.4.111/.119/.29/.58 + **§11.4.176** (the multi-track work-division + exactly-once-claim + device-lock arbitration anchor). Propagation gate `CM-COVENANT-114-178-PROPAGATION` (literal `11.4.178`) + recommended gate `CM-TRACK-QUALIFIED-IDENTITY` + paired §1.1 mutation. **Multi-project-per-track layout clarification (2026-07-10):** a `/mnt/track<N>` track parent is NOT single-project — MULTIPLE projects coexist on the same track parents, each in its OWN project-named subdirectory `/mnt/track<N>/<project_name>/` (e.g. `/mnt/track1/atmosphere/` AND `/mnt/track1/helix_ota/` side by side); the per-track working copy of project P on track N is ALWAYS `/mnt/track<N>/<project_name>/` (canonical lowercase snake_case dir-name per §11.4.29), never the bare track root and never a shared basename; this is WHY the identity carries the `<project>` component (`<project>__<track>__<role>`) — so two coexisting projects' track-N sessions / locks / logs / device-claims / registry rows never collide; a project adopting the multi-track engine (§11.4.187) MUST create AND register its own `/mnt/track<N>/<project_name>/` per track and MUST NOT assume it owns the track root or that a track hosts only one project. **Canonical authority:** constitution submodule `Constitution.md` §11.4.178. Non-compliance is a release blocker. No escape hatch — no `--bare-basename-ok`, `--skip-track-qualifier`, `--shared-session-name` flag.

§11.4.179 — Corruption-isolated parallel git streams (own-.git independent clones, not shared-common-dir worktrees) (User mandate, 2026-07-04). Compact summary: when corruption-isolation is a requirement, parallel git work streams MUST each be an isolated repo with its OWN `.git` (own object store, own index, own lock namespace), NOT `git worktree` checkouts sharing one common `.git`; a shared common-dir is a single point of failure — one stale `index.lock` / `.commit_all.lock` / `.push_all.lock` freezes EVERY stream at once (observed 9-h all-track freeze 2026-07-04) + one corrupted object store loses every stream; each stream's `git rev-parse --git-common-dir` MUST resolve INSIDE its own tree, a common-dir resolving to a shared parent is a violation; where CoW/reflink disk-feasibility is the constraint use the §11.4.167 reflink-clone-with-own-`.git` path (btrfs/XFS/ZFS/APFS reflink shares extents on disk while keeping a per-stream `.git`, so isolation + disk-feasibility both hold); honest boundary §11.4.6 — own-`.git` contains lock/corruption blast radius, it does NOT remove per-repo stale-lock reaping (§11.4.180) nor ff-only no-force integration (§11.4.113). Classification: universal (§11.4.17). Composes §11.4.167/.119/.58/§9.2 + **§11.4.176** (the multi-track work-division + exactly-once-claim + device-lock arbitration anchor). Propagation gate `CM-COVENANT-114-179-PROPAGATION` (literal `11.4.179`) + recommended gate `CM-GIT-STREAMS-OWN-COMMON-DIR` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.179. Non-compliance is a release blocker. No escape hatch — no `--shared-common-git`, `--worktree-when-isolation-required`, `--skip-git-isolation` flag.

§11.4.180 — Commit/push (single-writer) wrappers MUST auto-reap provably-stale locks before acquiring (User mandate, 2026-07-04). Compact summary: every commit / push / single-writer wrapper MUST, before acquiring its own lock, auto-reap any PROVABLY-stale git lock — one whose recorded holder PID is DEAD (`kill -0` fails), OR (no PID recorded) older than a defined threshold AND with no live wrapper/git process holding that repo; the reap covers the wrapper's own lock plus the git-internal existence-based locks (`index.lock`, `HEAD.lock`, `refs/**/*.lock`) git does NOT auto-release on holder death; it MUST NEVER remove a lock whose holder is ALIVE (removing a live-held lock corrupts a concurrent writer — §9.2); each reap decision (REAPED / KEPT + reason) is logged as captured evidence (§11.4.6 — liveness PROVEN via `kill -0`, never assumed); a wrapper that blocks indefinitely on a dead-holder lock (observed 9-h freeze 2026-07-04) is a violation; honest boundary §11.4.6 — reaping clears dead-holder deadlock, it does NOT substitute for the own-`.git` isolation of §11.4.179. Classification: universal (§11.4.17). Composes §11.4.84/.88/.6/§9.2/§11.4.116/§11.4.179. Propagation gate `CM-COVENANT-114-180-PROPAGATION` (literal `11.4.180`) + recommended gate `CM-WRAPPER-STALE-LOCK-AUTO-REAP` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.180. Non-compliance is a release blocker. No escape hatch — no `--skip-stale-lock-reap`, `--force-remove-live-lock`, `--block-on-dead-holder` flag.

§11.4.181 — Consistent controlled feature-branch naming: one feature/logic-group ⇒ exactly ONE canonical branch name across the main repo + all owned submodules (User mandate, 2026-07-05). Compact summary: one feature / logical group of workable items MUST map to EXACTLY ONE canonical branch name used IDENTICALLY on the main repo AND every touched owned submodule — NEVER two differently-named branches for one feature/group, NEVER per-repo naming drift; (1) single canonical name per group minted ONCE per the §11.4.167 D scheme (`feature/<slug>` branch + `<prefix>-<base>-feat-<slug>` tag, `<slug>` lowercase snake/kebab §11.4.29), used identically across the main repo + every branched submodule (an untouched submodule stays on its base branch); (2) single source of truth — the canonical name is recorded ONCE in the §11.4.176 exactly-once claim registry / the §11.4.93 workable-items `logic_group → branch` mapping and LOOKED UP, never re-invented (creating a branch whose name ≠ the group's registered canonical name, or minting a 2nd name for a group that already has one, is a §11.4.6 no-guessing violation); (3) mechanical enforcement (§11.4.75) — recommended gate `CM-BRANCH-NAME-CONSISTENCY` FAILs on an unregistered feature-branch name / two divergent names for one group / a submodule branch name diverging from the group's main-repo canonical name, paired §1.1 mutation (register one group under two divergent names → gate FAILs); (4) risk-free reconciliation (§9.2/§11.4.113/§11.4.84) — a mis-named/duplicate branch is reconciled by MERGING its commits onto the canonical branch (union, no commit lost, no `-s ours`/reset), ff-only push to all upstreams (NEVER force-push), retiring the mis-named branch only after the merge is confirmed on all mirrors. Classification: universal (§11.4.17). Composes §11.4.167/.176/.178/.179/.28/.29/.93/.113/.84/.6/.75/§9.2. Propagation gate `CM-COVENANT-114-181-PROPAGATION` (literal `11.4.181`) + recommended gate `CM-BRANCH-NAME-CONSISTENCY` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.181. Non-compliance is a release blocker. No escape hatch — no `--allow-divergent-branch-names`, `--skip-branch-registry`, `--rename-branch-freely`, `--per-repo-branch-name-OK`, `--two-names-one-feature-OK` flag.

§11.4.182 — Track+branch work-stream identity label `(T<N>/<branch>)` on every agent/subagent label + every operator-facing work-stream reference (User mandate, 2026-07-05). Compact summary: EVERY agent / subagent / work-stream label AND every operator-facing reference to a work stream MUST START with a `(T<N>/<branch> - <alias>)` prefix identifying the track + git branch + the claude-toolkit alias in use (e.g. `(T1/main - claude1) ATM-312 MPV drm_prime investigation`, `(T2/feature/mistiq-vader - deepseek) SPK-XXX ...`) so parallel-track work (`/mnt/track1..4`, each on its own branch, each driven by a distinct toolkit alias) is never ambiguous; (1) universal label form on every agent-dispatch description, status report, resume/handoff doc, progress line; (2) DETERMINISTIC derivation never guessed (§11.4.6) — `<N>` from the checkout path `/mnt/track<N>/…` (honest `?` off-track, never a guessed number), `<branch>` from `git rev-parse --abbrev-ref HEAD`, `<alias>` from `CLAUDE_CONFIG_DIR` basename `.claude-<alias>` → `<alias>` (honest `?` when unset/non-matching, never a guessed name) (reference labeler `scripts/multitrack/track_branch_label.sh`); (3) mechanical enforcement (§11.4.75 / §11.4.109-class) — a PreToolUse guard hook (`constitution/scripts/hooks/guard-track-branch-label.sh`, inherited BY REFERENCE per §11.4.177, never copied) BLOCKS any Agent/Task/TaskCreate dispatch whose `description` (or `subagent` label) does not START with `^\(T[0-9]+/[^)]+\)`, non-agent tools untouched; pre-build gate `CM-TRACK-BRANCH-LABEL` (labeler+hook exist/exec/parseable, hook wired in settings, convention doc exists) + propagation gate `CM-COVENANT-114-182-PROPAGATION` (literal `11.4.182`); (4) honest boundary (§11.4.6) — a `?` in the track OR the alias field is a valid honest label, never a bluff, and the enforced numeric-track format surfaces (blocks) an off-track dispatch rather than mislabeling it. Classification: universal (§11.4.17). Composes §11.4.178 (track-qualified identity — the label-prefix instantiation applied to every agent/subagent label + operator-facing reference) / §11.4.29 / §11.4.75 / §11.4.109 / §11.4.177 / §11.4.6 / §11.4.58 / §11.4.103 / §11.4.119 / §11.4.116 / §11.4.147. Propagation gate `CM-COVENANT-114-182-PROPAGATION` (literal `11.4.182`) + recommended gate `CM-TRACK-BRANCH-LABEL` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.182. Non-compliance is a release blocker. No escape hatch — no `--skip-track-label`, `--unlabeled-agent-OK`, `--bare-work-stream-reference`, `--no-track-branch-prefix`, `--label-from-memory` flag.

§11.4.183 — Maximal multi-agent utilization per work-stream + full-constitution-application + zero-bluff mandate (User mandate, 2026-07-07). Compact summary: every track / work-stream MUST maximize multi-agent working approaches — subagent-driven development (§11.4.20/§11.4.70), independent review agents (§11.4.125/§11.4.142/§11.4.165/§11.4.134), parallel background streams with auto-backfill (§11.4.58/§11.4.103), and every other applicable agent form — so every track produces the MOST completed work at the HIGHEST quality in the LEAST time; every track MUST apply the ENTIRE constitution (every rule, mandatory constraint, guide, and derived technology — NOTHING ignored, skipped, avoided, or deemed less-relevant); there MUST NEVER be false / faulty / untested / unverified results or bluff of any kind anywhere (§11.4 anti-bluff covenant, captured evidence §11.4.5/§11.4.69/§11.4.107); honest boundary §11.4.6 — "maximize agents" is bounded by §12.6 60% memory + §11.4.58 parallel-agent cap + genuine parallelizability + §11.4.119 single-resource-owner, spawning contending/redundant agents is NOT maximization, the bar is maximum USEFUL parallel throughput not agent count. Classification: universal (§11.4.17). Composes §11.4.20/.58/.70/.92/.94/.97/.103/.119/.125/.126/.134/.142/.165/§12.6. Propagation gate `CM-COVENANT-114-183-PROPAGATION` (literal `11.4.183`) + recommended gate `CM-MAX-AGENTS-PER-TRACK` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.183. Non-compliance is a release blocker. No escape hatch — no `--single-agent-OK`, `--skip-review-agents`, `--partial-constitution-application`, `--serialise-parallelisable-work` flag.

§11.4.184 — Mandatory SonarQube static-analysis CLI + local tooling installed and PATH-discoverable (User mandate, 2026-07-06). Compact summary: every project/host that runs static analysis MUST have the SonarQube scanner CLI (`sonar-scanner`) installed AND durably discoverable on the system PATH via the shell rc (`.bashrc` / `.zshrc`) — exactly like the Firebase/GitHub/GitLab CLIs — PLUS the shared `constitution/scripts/sonarqube/` tooling (`sonarqube_install_check.sh` / `sonarqube_lib.sh` / `sonarqube_run_scan.sh` / `sonarqube_container.sh` + `compose/`) consumed BY REFERENCE (§11.4.28 decoupled + §11.4.177 project-agnostic + §11.4.80-style inherited, NEVER copied per project); (1) installed + PATH-durable — the scanner resolves in every fresh shell via a durable rc export, not an ephemeral session-only PATH; (2) shared tooling present + GREEN — `sonarqube_install_check.sh` exits 0 (scanner + rootless podman §11.4.161 + podman-compose + curl + ES-ready `vm.max_map_count`) before any scan-dependent work; (3) anti-bluff §11.4.6 — "installed/configured" PROVEN by `command -v sonar-scanner` resolving AND install-check exit 0, NEVER assumed, a claimed-installed-but-absent CLI is a §11.4 tooling-layer bluff; (4) rootless §11.4.161 — the local SonarQube server runs via rootless podman, never rootful docker/sudo; honest boundary §11.4.6 — names the SonarQube SCANNER CLI + shared local-server tooling, and per §11.4.166 (former universal Semgrep static-analysis mandate REPEALED) SonarQube is a DISTINCT operator-mandated tool not a Semgrep re-instatement, mandating TOOLING availability not a scan on every build. Classification: universal (§11.4.17). Composes §11.4.28/.36/.74/.75/.109/.161/.166/.177/.6. Propagation gate `CM-COVENANT-114-184-PROPAGATION` (literal `11.4.184`) + recommended gate `CM-SONARQUBE-CLI-INSTALLED` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.184. Non-compliance is a release blocker. No escape hatch — no `--skip-sonarqube-install`, `--sonar-cli-optional`, `--ad-hoc-sonar-path`, `--rootful-sonar-OK` flag.

§12.12 — Process/thread-limit (RLIMIT_NPROC) awareness for parallel subagent/multi-process work (User mandate, 2026-07-07). Compact summary: heavy parallel subagent/multi-process work is bounded by the OS per-user process/thread limit `ulimit -u` (`RLIMIT_NPROC`, Linux counts ALL threads+processes of the real UID, not per-process); the ambient tooling fleet (MCP servers, tokio/V8/libuv worker runtimes, `mongod` / `chrome` / `prisma`-class MCP servers, local LLM servers) consumes a large FIXED fraction before project work starts — FORENSIC FACT (2026-07-07): a host at `ulimit -u 4096` sat at ~3900-4005 threads from the ambient fleet alone (<150 headroom); a 4th Go-heavy subagent (`GOMAXPROCS=nproc` spawns dozens of threads) hit `EAGAIN` / `failed to create new OS thread (errno=11)` and crashed tooling; stopping 3 non-essential MCP servers freed ~1345 threads (3827→2482). Mandate: BEFORE scaling parallel subagents/processes, check thread headroom (`ulimit -u` soft + `ulimit -Hu` hard + live `ps -L --no-headers -u "$USER" | wc -l`) and treat exhaustion as a §12 host-safety event that YIELDS UNCONDITIONALLY — the ORTHOGONAL second exhaustion axis alongside §12.6's memory ceiling (abundant RAM ≠ abundant thread budget); signals `EAGAIN` / `fork: retry` /

`failed to create new OS thread` / `cannot allocate memory` on fork/clone (distinct from OOM, §11.4.6 no-guessing — never conflate); low headroom ⇒ SERIALIZE, avoid `GOMAXPROCS=nproc`-class spikes, never dispatch blind (a subagent hitting `EAGAIN` mid-`git push` is a §9.2 data-safety risk). Three raising techniques: (a) no-restart `prlimit --pid <PID> --nproc=<soft>:<hard>` (root, e.g. `su -c`) — new children of a running process inherit the raise; (b) persistent `/etc/security/limits.conf` (`<user> hard/soft nproc <N>`) or systemd `LimitNPROC=<N>`, next login; (c) self-raise `ulimit -u <N>` up to the existing hard ceiling within the current session — `su -c "ulimit -u N"` sets the limit ONLY inside that transient subshell, NOT the already-running session. Freeing threads (stopping non-essential MCP servers) REQUIRES operator authorization (§11.4.122, no autonomous tooling removal); CAUTION `pkill -f` / `pgrep -f` match the killer's OWN command line — always exclude the caller's `$$` / `$PPID` (and conductor PID) from any cmdline-pattern kill list to avoid self-kill. Classification: universal (§11.4.17). Composes §12/.6/.7/§11.4.58/.101/.103/.122/§9.2. Propagation gate `CM-COVENANT-12-12-PROPAGATION` (literal `12.12`) + recommended gate `CM-NPROC-HEADROOM-CHECK` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §12.12. Non-compliance is a release blocker. No escape hatch — no `--ignore-thread-limit`, `--dispatch-blind`, `--skip-headroom-check`, `--autonomous-tooling-kill`, `--assume-ulimit-scope` flag.

§11.4.185 — Manual QA-team testing as the mandatory FINAL confirmation of done (User mandate, 2026-07-07). Compact summary: no set of features / scope of work / release may be considered fully completed until it has received MANUAL TESTING confirmation by the project's QA team as the FINAL confirmation step — everywhere and always, never forgotten; every automated gate (§11.4 anti-bluff covenant, §11.4.5/§11.4.69 captured evidence, §11.4.52 autonomous validation, §11.4.40 full-suite retest, §11.4.108 four-layer verification, §11.4.132 risk-ordered validation) is NECESSARY but NOT SUFFICIENT — manual QA is the FINAL human sufficiency gate; ADDS to (never weakens) §11.4.52's mandatory autonomous-validation-path requirement — a feature needs BOTH the autonomous path AND the manual-QA final confirmation, never one substituting the other; pipeline: autonomous-GREEN → build → flash/deploy → autonomous on-device validation → hand off to QA team → ONLY after QA-team manual confirmation is the scope fully-completed / tag-eligible; the §11.4.126 autonomous-loop release-tag terminal condition now REQUIRES this confirmation; honest boundary §11.4.6 —

the agent MUST hand off + WAIT, never self-certify/simulate/infer the manual step, while continuing every other non-blocked item in parallel per §11.4.87/.94/.97/.101 (the wait parks one scope, not the loop). Classification: universal (§11.4.17). Composes §11.4/.5/.6/.40/.52/.69/.108/.126/.129/.130/.132. Propagation gate `CM-COVENANT-114-185-PROPAGATION` (literal `11.4.185`) + recommended gate `CM-MANUAL-QA-FINAL-CONFIRMATION` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.185. Non-compliance is a release blocker. No escape hatch — no `--skip-manual-qa`, `--automation-suffices`, `--self-certify-manual-step`, `--assume-qa-confirmed`, `--autonomous-only-completion` flag.

§11.4.186 — Anti-divergence enforcement: cross-document consistency is a mandatory-before-export/commit gate, never an after-the-fact audit (research-derived, 2026-07-08). Compact summary: any project maintaining more than one representation of the same tracked data (a delivery plan + its MVP brief + a workable-items DB + their summaries + `.md` / `.html` / `.pdf` / `.docx` exports) MUST enforce cross-document CONSISTENCY as a deterministic gate that runs BEFORE any export render, any doc/DB sync `verify`, and any doc-set commit — NEVER as an after-the-fact human audit that runs only when an operator happens to ask (the forensic FACT: the same NanoKVM/phone-as-screen feature appeared in three plan clusters — one shared ticket SPK-596 across SIX rows split over TWO releases, and a carve-out due to ship 2026-08-31 whose own prerequisite does not START until 2026-10-05 — and NO gate caught it; only an operator's manual read of `Plan_v9.0.xlsx` surfaced it, after the divergent bytes had already exported). (1) **One no-divergence verdict, three seams** — a single deterministic PASS/FAIL/SKIP verdict (built by composing a cross-document integrity validator with the workable-items-DB internal `validate`) gates all three write-seams (export / sync-verify / commit); a FAIL at any seam refuses that seam, naming the exact offending records (§11.4.6 actionable). (2) **Five decidable check families** — DEDUP (no two records describe one feature with divergent timeline/status/release, keyed on ticket OR normalised `(subject, scope)`, NOT bare subject substring), TIMELINE (start ≤ deadline; no deadline-before-dependency; no dependency cycle; GATED exempt-but-marked), CROSS-DOC (the same item across every representation carries identical timeline/status/type against a NAMED authoritative source), INTEGRITY (no orphan refs; Status↔Type §11.4.33; location↔status), STRUCTURAL (required columns; ID uniqueness+monotonicity §11.4.54). (3) **Drift-proof trigger** — a §11.4.86 sha256-of-sorted-members fingerprint of the authoritative inputs re-arms the gate on ANY input change, so a stale verify cannot pass on changed data. (4) **Self-validated analyzer (§11.4.107(10))** — the validator ships golden-good (MUST PASS) + one golden-bad per family (MUST FAIL, pinpointing the offender) + a negative-control (distinct same-subject tasks that MUST PASS — the false-positive guard); a validator that passes its golden-bad, or fails golden-good/the negative-control, is itself a §11.4 bluff and a release blocker. (5) **Reusable + decoupled (§11.4.28/.31)** — packaged as a depth-1 reusable engine under the constitution submodule per the §11.4.28(C) carve-out with a `helix-deps.yaml`; project-specific doc-sets + authoritative-source bindings live in a consumer-owned checkset (data, never engine code). (6) **Supersedes ad-hoc audits** — the per-cycle `*_INTEGRITY_FINDINGS.md` / `RECONCILIATION_*.md` after-the-fact audit docs are retired in favour of the gate; honest boundary §11.4.6 — the gate proves internal CONSISTENCY (no record diverges, no timeline is incoherent, every ref resolves), NOT plan CORRECTNESS (achievability/date-rightness stay operator/management decisions the gate SURFACES, never MAKES). Classification: universal (§11.4.17). Composes §11.4.6/.12/.33/.44/.50/.53/.54/.60/.65/.73/.75/.85/.86/.91/.93/.95/.106/.107/.108/.110/.148/.176. Propagation gate `CM-COVENANT-114-186-PROPAGATION` (literal `11.4.186`) + recommended gate `CM-DOC-INTEGRITY-VALIDATION` (5 invariants: validator present+executable; consumer checkset present; the no-divergence gate wired into each of the three seams; fingerprint sidecar fresh; `selfcheck` golden-good/golden-bad/negative-control wired into meta-test) + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.186. Non-compliance is a release blocker. No escape hatch — no `--skip-divergence-gate`, `--audit-after-export-OK`, `--cross-doc-check-not-applicable`, `--unvalidated-analyzer-OK`, `--divergent-commit-OK` flag.

§11.4.187 — Automatic multi-track ruler orchestration (User mandate, 2026-07-04; landed 2026-07-09). Compact summary: every project incorporating this constitution MUST ship its multi-track parallel-development orchestration as a UNIVERSAL, AUTOMATIC, OUT-OF-THE-BOX, INHERITED capability — not a per-project reimplementation — via a single conductor session that programmatically SPAWNS, DRIVES, RESUMES, MONITORS, and CONTROLS per-track headless workers. (1) **Headless spawn primitive** — workers launch via native headless `claude -p --output-format stream-json --verbose`; the worker's `session_id` is captured from the FIRST `init` event (never invented);

success/failure is read from the terminal event's `result.is_error` field, NEVER the process exit code (a CLI can exit non-zero on a successful `result`, or zero on a failed one — exit-code-as-success-signal is a §11.4.6 guess); continuation of an existing worker uses `--resume <session_id>` invoked from the SAME cwd + env the spawn used (the storage/resume-collision gotcha). (2) **Per-subscription auth + rate-limit fallback** — every spawn `unset`s any ambient `ANTHROPIC_API_KEY` (the precedence trap that silently routes a "subscription" worker onto pay-per-token API billing) and sets the selected alias's `CLAUDE_CODE_OAUTH_TOKEN` / `CLAUDE_CONFIG_DIR`; on a captured rate-limit signature (`isApiErrorMessage` + `apiErrorStatus:429`, quota-exhausted vs transient distinguished by evidence, never guessed) the ruler REBINDS the affected track onto the next healthy alias (cross-track reuse permitted), falls back to a provider alias where configured, and bounded-parks (§11.4.101) rather than hanging when every alias is cooled. (3) **Crash-resilient ruler self-supervisor** — the ruler's OWN state (per-track alias/session_id/worktree/last-verdict) is persisted durably (write-temp-then-rename, NEVER tmpfs) so an external watchdog detects a dead ruler PID and re-hydrates every track's binding from the durable state + the §11.4.116 event stream — no work is silently lost to a ruler crash (§11.4.147 at the ruler layer). (4) **Idempotent bootstrap + auto-install + conductor-stays-home** — one idempotent bootstrap command wires the whole system out of the box (cwd-hook install, config validation, reconciliation), auto-invoked on every constitution pull via the §11.4.164 `post_update_hook.sh` skill-register seam; the `conductor:` -designated alias/session is NEVER bound into a track worktree (the resolver returns nothing for it) so the conductor always runs from the main checkout. (5) **Project-agnostic engine, consumer-owned data** — the entire mechanism (config loader, alias→worktree resolver, cwd-hook, bind/fallback/cooldown orchestrator, session launcher, rate-limit monitor, supervisor) lives at `constitution/scripts/multitrack/`, inherited BY REFERENCE (§11.4.28(B)/§11.4.177), carrying NO project literal; every consuming project supplies its own per-host `config/multitrack/<hostname>.yaml` (tracks/mounts/serials/aliases/conductor) as DATA, never as an engine edit. Anti-bluff (§11.4/§11.4.108): every piece carries four-layer coverage (pre-build gate + on-host test + paired §1.1 mutation + captured evidence per §11.4.69), and each fix declares a runtime signature verified on a clean host — never a claim the mechanism "should work." Honest boundary (§11.4.6): this anchor mandates the MECHANISM; genuine per-subscription quota-isolation proof requires live tokens exercising real rate-limit boundaries (an operator-scheduled live acceptance window) — a mock-driven GREEN proves the spawn/resume/fallback CONTRACT, never live-quota behaviour, and the two are never conflated or substituted for each other. Classification: universal (§11.4.17). Composes §11.4.20 (subagent-driven-by-default) / §11.4.28 (submodule decoupling + dependency layout) / §11.4.58 (parallel-development PWU pipeline) / §11.4.70 (subagent-driven execution default) / §11.4.101 (autonomous-decision-over-blocking — the bounded-park-on-all-cooled behaviour) / §11.4.103 (continuous parallel-stream routine) / §11.4.111 (resolve-by-stable-name — config by serial, never by index) / §11.4.116 (real-time conductor↔framework sync channel) / §11.4.119 (single-resource-owner partitioning) / §11.4.126 (default autonomous-loop mode) / §11.4.128 (always-on device-recording — the analogous host-safety discipline) / §11.4.131 (standing session-resumption file) / §11.4.147 (crashed-agent respawn-until-complete — the ruler-self-supervisor is its ruler-layer instantiation) / §11.4.164 (constitution auto-propagation hook seam) / §11.4.167 (feature work-stream lifecycle) / §11.4.176 (multi-track work-division + exactly-once claim + device-lock) / §12.6/§12.8 (host-safety memory + concurrency ceilings — the spawn/launch pool is budget-gated). Propagation gate `CM-COVENANT-114-187-PROPAGATION` (literal `11.4.187`) + recommended gates `CM-MULTITRACK-ENGINE-IN-CONSTITUTION` (the generic ruler scripts present under `constitution/scripts/multitrack/`, carrying no project literal, `bash -n` + `sh -n` clean per §11.4.67) / `CM-MULTITRACK-AUTOINSTALL-WIRED` / `CM-MULTITRACK-BOOTSTRAP-IDEMPOTENT` / `CM-MULTITRACK-BIND-ON-START` / `CM-MULTITRACK-FALLBACK-MONITOR` / `CM-MULTITRACK-HEADLESS-SPAWN` / `CM-MULTITRACK-RULER-SUPERVISOR` + paired §1.1 mutation (strip the literal → propagation gate FAILs; re-introduce a project literal into the engine, or downgrade a rate-limit rebind to exit-code-based success detection → the corresponding mechanism gate FAILs). **Canonical authority:** constitution submodule `Constitution.md` §11.4.187. Non-compliance is a release blocker. No escape hatch — no `--skip-multitrack`, `--manual-multitrack-only`, `--no-auto-install`, `--per-project-reimpl-OK`, `--exit-code-is-success-signal` flag.

§11.4.188 — Regular main→feature merge cadence: every track / feature-branch MUST FREQUENTLY merge canonical `main` into itself DURING the work, never only at the end (User mandate, 2026-07-09). Compact summary: every long-lived feature branch AND every parallel-development track (§11.4.58/§11.4.103 parallel streams; §11.4.167 feature work-streams; §11.4.176/§11.4.178/

§11.4.179 multi-track) MUST regularly `git merge origin/main` (the canonical trunk/master) INTO its own branch THROUGHOUT the work — NOT only at the end — so no branch ever drifts far from main and the eventual back-merge stays a small, low-conflict, low-risk operation. GENERALISES §11.4.167(D) (trunk-merged-INTO-every-BIG-feature-STREAM, at minimum after every trunk tag / daily, `git merge` never rebase, stale-if-not-synced) from the big-work-item feature-stream lifecycle to EVERY feature branch on EVERY track — a small/medium feature branch, or any `git worktree` /CoW track that never earned a full §11.4.167 stream, is STILL bound to this cadence (the §11.4.167(D) discipline promoted to the general standalone rule; §11.4.167(D) remains its big-item specialisation). (1) **WHAT** — fetch the latest canonical `main` / `master` and merge it INTO the working branch, keeping the branch continuously close to trunk; a branch that has NOT integrated trunk within the cadence is flagged STALE (the exact long-lived-branch conflict-storm + integration-risk anti-pattern this forbids). (2) **CADENCE** — merge trunk in FREQUENTLY: after EVERY trunk tag, at minimum DAILY while the branch is live, AND before starting any significant new chunk of work on the branch — MANY SMALL merges, never one big deferred end-of-branch merge (small frequent merges minimise the divergence + conflict surface). (3) **HOW (safe-by-construction)** — MERGE, NEVER rebase a shared/tagged/pushed branch (no commit-loss, no history rewrite — §9/§11.4.113); FETCH-FIRST every remote before the merge (§11.4.37/§11.4.71 — remote state is unknowable without a fetch, §11.4.6); take a §9.2 hardlinked pre-op backup before any LARGE / conflict-heavy / risky merge; resolve every conflict CAREFULLY — ZERO conflict markers committed, ZERO file silently dropped, union of both sides preserved (§11.4.41 step 3 / §9 no-loss); integrate ONLY when the branch is QUIESCENT (§11.4.84 — no in-flight mutation gate, no unaccounted uncommitted work); prefer running the merge + its verification in the BACKGROUND parallel to the main stream so it never stalls other work (§11.4.88/§11.4.89/§11.4.103); NEVER force-push the merged result (§11.4.113 — the merged branch fast-forwards its own mirrors). (4) **CONSISTENCY + ANTI-BLUFF** — after EVERY merge the branch MUST be PROVEN consistent, never ASSUMED: a pre-build smoke gate runs post-merge (§11.4.42 step-3 smoke) and is GREEN, a `git grep` for conflict markers (`<<<<<<<` / `=======` / `>>>>>>>`) returns EMPTY, and no commit present pre-merge on either side is lost — captured evidence per §11.4.5/§11.4.69; a merge claimed "clean" without the post-merge smoke + marker-scan + no-lost-commit check is a §11.4/§11.4.1 bluff at the integration layer. (5) **HONEST BOUNDARY (§11.4.6)** — regular trunk-integration keeps the branch MERGEABLE + low-conflict;

it does NOT by itself prove the branch's own work is correct (each change still crosses §11.4.108 runtime-signature + §11.4.40 retest), and it is NOT the approval-gated back-merge TO trunk (that stays §11.4.167(I)/§11.4.40/§11.4.41 no-merge-until-approved). Classification: universal (§11.4.17) — the consuming project supplies its canonical-trunk ref name, cadence unit, post-merge smoke-gate command, and per-track worktree/CoW mechanism per §11.4.35. Composes §9/§9.2/§11.4.6/§11.4.17/§11.4.37/§11.4.41/§11.4.42/§11.4.71/§11.4.84/§11.4.88/§11.4.89/§11.4.103/§11.4.113/§11.4.167/§11.4.176/§11.4.178/§11.4.179/§11.4.181. Propagation gate `CM-COVENANT-114-188-PROPAGATION` (literal `11.4.188`) + recommended gate `CM-REGULAR-MAIN-MERGE-INTO-FEATURE` (every live feature branch / track has integrated the canonical trunk within the declared cadence; its last integration is a real `git merge` not a rebase; the post-merge smoke + zero-conflict-marker + no-lost-commit checks passed) + paired §1.1 mutation (let a live branch go stale past the cadence, OR record a rebase of a shared branch in place of a merge, OR commit a conflict marker → the gate FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.188. Non-compliance is a release blocker. No escape hatch — no `--skip-main-merge`, `--merge-only-at-end`, `--rebase-shared-branch`, `--defer-trunk-sync`, `--no-post-merge-smoke`, `--force-push-after-merge`, `--big-items-only-trunk-sync` flag.

§11.4.189 — Most-reopened cases get extra-depth live-testing scrutiny FIRST (User mandate, 2026-07-10). Verbatim operator mandate: "All live testing MUST pay extra attention to cases which have been reopened numerous times and those cases in depth retested, investigated, fully validated and verified with real physical results and no bluff of any kind anywhere!" All LIVE TESTING (on-device / on-target runtime validation, deploy-and-observe) MUST give EXTRA-DEPTH retest + in-depth investigation + FULL validation/verification with REAL PHYSICAL captured evidence (§11.4.5/§11.4.69/§11.4.107 — captured audio/video/sysfs/dumpsys/sink-side/runtime-signature) and NO bluff of any kind ANYWHERE to the cases that have been REOPENED THE MOST TIMES (highest §11.4.55 reopens-count) — the empirically-most-fragile set gets the DEEPEST live scrutiny, and it gets it FIRST. (1) **WHAT** — for the highest-reopens-count cases, live testing is NOT a single pass: it re-runs the case at EXTRA DEPTH (more iterations per §11.4.50 deterministic consistency, wider edge-case + regression coverage), RE-INVESTIGATES the underlying defect per §11.4.102 systematic-debugging + §11.4.146 reproduce-first, and CONFIRMS with rock-solid captured physical evidence per §11.4.123 — never metadata-only / config-only / absence-of-error / grep-without-runtime (§11.4/ §11.4.1), never a false result. (2) **ORDER** — the most-reopened set is scrutinised FIRST, ahead of the rest of the live-test suite (the §11.4.132 risk-descending discipline, with most-reopened the load-bearing key). (3) **KEY** — priority is keyed off the §11.4.55 reopens-count (the consuming project supplies its reopens-count source per §11.4.35, e.g. the §11.4.93 workable-items DB `reopens_count`); a high reopens-count is the strongest empirical fragility signal (§11.4.132(d)). (4) **PROOF** — each most-reopened case's live PASS cites its captured-evidence artefact path (§11.4.69/§11.4.107) on a CLEAN/fresh deployment (§11.4.108 runtime-signature), with the §11.4.115 RED→GREEN polarity flip captured where a permanent regression guard exists (§11.4.135). STRENGTHENS/REFINES §11.4.132 (risk-ordered validation factor (d) most-reopened — §11.4.189 promotes most-reopened into its OWN EXTRA-DEPTH mandate at the LIVE-TEST layer, not merely one ranking factor among four) + §11.4.55 (reopens-count as the priority key) + §11.4.129/§11.4.130 (validate-the-just-fixed / highest-risk-first after redeploy). Classification: universal (§11.4.17) — references no project-specific hardware; the consuming project supplies its reopens-count source per §11.4.35. Composes §11.4.5/§11.4.6/§11.4.55/§11.4.69/§11.4.107/§11.4.108/§11.4.115/ §11.4.129/§11.4.130/§11.4.132/§11.4.135/§11.4.146. Propagation gate `CM-COVENANT-114-189-PROPAGATION` (literal `11.4.189`) + recommended gate `CM-MOST-REOPENED-EXTRA-DEPTH-LIVE-TEST` (the live-test suite orders + extra-depth-scrutinises the highest-reopens-count cases FIRST, each with captured physical evidence on a clean deployment) + paired §1.1 mutation (drop the most-reopened extra-depth pass, OR run it AFTER the rest of the suite, OR PASS a most-reopened case on metadata-only evidence → the gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.189. Non-compliance is a release blocker. No escape hatch — no `--skip-most-reopened-depth`, `--reopened-same-as-any`, `--live-test-any-order`, `--metadata-pass-for-reopened-OK`, `--defer-reopened-scrutiny` flag.

§11.4.190 — Website engineering-quality mandate: every project website MUST be fully responsive + completely SEO-optimized + uniquely OpenDesign-authored + bleeding-edge enterprise-quality, each PROVEN with captured evidence (User mandate, 2026-07-10). Verbatim operator mandate: "every website we create inside our projects MUST BE fully responsive - working on all browsers, all operating systems, all platforms and all devices and screen sizes! Every Website MUST BE completely SEO optimized! Layouts and templates done MUST BE unique and everything fully designed from OpenDesign! Make sure this applies to our project Website as well! Everything must be bleeding-edge enterprise quality, stunning visual look and feel, innovative, creative and unique! There MUST BE no false or faulty result or bluff of any kind anywhere!" Every website / web-UI surface a project ships — INCLUDING this project's OWN website — MUST satisfy AND PROVE with captured physical evidence (§11.4.5/§11.4.69/ §11.4.107/§11.4.170 — never a bare claim, §11.4.6) ALL of: **(A) FULL RESPONSIVENESS** — correct, usable rendering across ALL major browser engines (Chromium/Firefox/WebKit), all operating systems, all device classes (desktop/laptop/tablet/phone/large-display) and the full range of screen sizes/ orientations/pixel-densities; NO broken layout / overflow / clipping / overlap / off-screen or unusable control at ANY breakpoint. **(B) COMPLETE SEO OPTIMIZATION** — semantic HTML + correct document outline; per-page unique `<title>` + meta description; Open Graph + Twitter cards; canonical URLs; structured data (schema.org / JSON-LD); `robots` + XML sitemap; crawlable / indexable markup; WCAG AA accessibility + Core Web Vitals performance (both feed ranking). **(C) UNIQUE OpenDesign-AUTHORED templates/layouts** — EVERY layout + template designed FROM the OpenDesign design-token system (§11.4.162), NOT generic / off-the-shelf / templated; visually unique, innovative, creative, brand-distinct. **(D) BLEEDING-EDGE ENTERPRISE VISUAL QUALITY** — stunning, polished, professional look-and-feel in BOTH light AND dark themes (§11.4.162 tokens). **(E) ANTI-BLUFF PROOF** — NONE of (A)–(D) may be CLAIMED without captured physical evidence: responsiveness PROVEN by device-independent host-rendered screenshots across the real breakpoint × engine matrix + an OCR/vision layout oracle (no overlap / label-over-label / clip / overflow / off-screen); SEO PROVEN by an automated audit (Lighthouse SEO + structured-data validation + meta / OG / sitemap / robots presence) meeting a defined score floor; visual quality PROVEN by the §11.4.170 device-independent host-rendered pixel proof per screen × state × {light,dark}; uniqueness PROVEN by OpenDesign-token provenance. A website reported responsive / SEO-optimized / unique / enterprise-quality WITHOUT these captured proofs is a §11.4 PASS-bluff at the web-UI layer. §11.4.190 is the strict WEB-quality expansion of §11.4.162 (OpenDesign — it applies §11.4.162's design-token mandate to the full responsive + SEO + enterprise-quality web surface) and binds §11.4.170's device-independent host-rendered pixel proof as the correctness oracle; sibling of §11.4.168 (exported-document independent visual validation) at the web-UI layer. Classification:

universal (§11.4.17) — references no project-specific hardware; the consuming project supplies its concrete breakpoint × engine matrix, SEO score-floor, and OpenDesign token set per §11.4.35. Composes §11.4.5/§11.4.6/§11.4.69/§11.4.107/§11.4.162/§11.4.168/§11.4.170. Propagation gate `CM-COVENANT-114-190-PROPAGATION` (literal `11.4.190`) + recommended gates `CM-WEBSITE-RESPONSIVE-PROVEN` (host-rendered screenshots across the breakpoint × engine matrix + layout oracle — no overlap/clip/overflow/off-screen) / `CM-WEBSITE-SEO-OPTIMIZED` (automated SEO audit meets the score floor + structured-data / meta / OG / sitemap / robots present) / `CM-WEBSITE-OPENDESIGN-UNIQUE` (every layout/template's OpenDesign-token provenance, no generic/off-the-shelf template) + paired §1.1 mutation (report a website responsive/SEO/unique/enterprise-quality with NO captured proof, OR ship a generic non-OpenDesign template, OR skip the light/dark host-rendered pixel proof → the recommended gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.190. Non-compliance is a release blocker. No escape hatch — no `--skip-responsive-proof`, `--seo-optional`, `--generic-template-OK`, `--visual-quality-later`, `--claim-without-render` flag.

§11.4.191 — Work-to-track/branch binding enforcement: a logic-group's work can NEVER be committed OR dispatched onto the wrong track/branch (research-derived, 2026-07-10). Compact summary: every logical group of workable items binds to EXACTLY ONE canonical `(track, branch)` recorded ONCE in the single source of truth (the §11.4.93 workable-items DB `logic_groups.destination` = authoritative branch + a `canonical_track` column + a `group_paths` file-scope manifest), and NO change belonging to a group may LAND on, nor be DISPATCHED to, a track/branch other than that group's canonical one — the file-scope generalisation of §11.4.181/§11.4.182 from branch-NAME consistency to work-PLACEMENT consistency (forensic FACT 2026-07-10: scrcpy/Remote-app work — group `mistiq-vader`, canonical `feature/mistiq-vader`/T2 — was committed onto `main`/T1; §11.4.181's `guard-branch-consistency.sh` never fired because committing onto an existing branch is not a feature-branch CREATE, and no gate checked which FILES land on which BRANCH). (1) **Single source of truth** — the `logic_group → (branch, track, path-globs)` binding is recorded ONCE in the DB and LOOKED UP, never re-invented (a commit/dispatch placing a group's file-scope on a non-canonical branch/track is a §11.4.6 no-guessing violation). (2) **Preventive hook**

(§11.4.109/§11.4.177) — a PreToolUse guard `guard-work-track-binding.sh` (inherited BY REFERENCE, never copied) BLOCKS (exit 2) a git-commit-class Bash call whose to-be-committed file set contains a path owned by a group whose canonical branch/track ≠ the current checkout, AND BLOCKS an Agent/Task dispatch whose §11.4.182 label track/branch ≠ the canonical (track, branch) of the ticket's group; fail-closed on unreadable registry, `# guardrails:allow <reason>` audited escape. (3) **Detective gate (§11.4.110/§11.4.108)** — `CM-WORK-TRACK-BINDING-ENFORCED` asserts, on the current branch's own added commits (`merge-base(HEAD,main)..HEAD`), that every `group_paths` -owned file's group destination == current branch (and canonical_track == current track when pinned) — the actual enforcement boundary regardless of how the commit was made, the runtime signature being FAIL-on-wrong-placement / PASS-on-canonical-placement on a clean checkout. (4) **Risk-free reconciliation (§9.2/§11.4.113/§11.4.84)** — a mis-placed change is reconciled by MERGING its commits onto the group's canonical branch (union, no commit lost, no `-s ours` /reset), ff-only push to all upstreams (NEVER force-push), then removing it from the wrong branch only after the merge is confirmed on all mirrors. Classification: universal (§11.4.17) — references no project hardware; the consuming project supplies its `(track, branch)` map, `group_paths` seed, and commit-wrapper wiring per §11.4.35. Composes §11.4.6/.28/.29/.35/.58/.75/.84/.93/.95/.109/.110/.113/.119/.167/.176/.177/.178/.181/.182/ §9.2. Propagation gate `CM-COVENANT-114-191-PROPAGATION` (literal `11.4.191`) + recommended gate `CM-WORK-TRACK-BINDING-ENFORCED` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.191. Non-compliance is a release blocker. No escape hatch — no `--allow-cross-track-commit` , `--skip-work-binding` , `--commit-any-branch-OK` , `--dispatch-any-track-OK` , `--no-file-scope-check` flag.

§11.4.192 — Continuous multi-track auto-backfill: a FREE track MUST be IMMEDIATELY re-assigned its next-highest-priority domain work-set, never left idle while its domain has actionable items (User mandate, 2026-07-11). Verbatim operator mandate: "once a Track is free, we shall assign next set of workable items to it ... full continuous multi-track development ... no stopping ... this is MANDATORY". Compact summary: the moment any track's worker COMPLETES its current work-set (its item(s) closed with anti-bluff evidence, or merged), that FREE track MUST be IMMEDIATELY assigned the NEXT-highest-priority actionable work-set from its OWN domain (its §11.4.191 logic_group /

canonical branch) — a track is NEVER left idle while its domain still holds actionable workable items, and the conductor/ruler CONTINUOUSLY re-assigns with NO stopping (full continuous multi-track development); PROMOTES the §11.4.103(B) ≥3-parallel-background-streams-with-auto-backfill invariant into the multi-track (§11.4.176/§11.4.178/§11.4.179/§11.4.187) layer (§11.4.103(B) backfills a STREAM the moment it is done; §11.4.192 backfills a TRACK the moment its worker is done, keyed to that track's canonical (track, branch) domain per §11.4.191). (1) **WHAT** — on any track-worker completion the conductor/ruler (§11.4.187) MUST, per that track's domain, look up the NEXT actionable work-set, CLAIM it exactly-once (§11.4.176 exactly-once claim registry, no two tracks claiming one item), and DISPATCH it as the track's new §11.4.182-labelled work stream — without waiting for a scheduled wake or an operator prompt. (2) **PRIORITY** — the next work-set is the track domain's highest-priority actionable item (§11.4.42/§11.4.72 priority order, audio-first where in scope), never an arbitrary pick. (3) **SOURCE-SIDE ADVANCES EVEN WHEN VALIDATION IS GATED** — when the next item's on-device VALIDATION is blocked on a pending build / device availability (§11.4.108 clean-target / §11.4.185 manual-QA / operator-gated flash), the track MUST still advance the item's SOURCE-SIDE work (investigation, RED-test §11.4.115, source patch, pre-build gate + paired §1.1 mutation, docs) and honestly §11.4.3/§11.4.52 SKIP-with-reason ONLY the genuinely device-gated validation portion — a build/device gate on the validation step is NOT a reason to idle the whole track. (4) **HONEST BOUNDARY (§11.4.6)** — a track idles ONLY when its domain GENUINELY has zero remaining actionable workable items OR every remaining item is externally blocked with no source-side advancement possible (§11.4.94(A)/§11.4.97/§11.4.101 idle-only-when-blocked closed set); "the merges are done" / "the current item is done" / "nothing visible right now" is NEVER a valid idle reason while the domain still holds progressable items (a §11.4.94+§11.4.97 violation); an unavoidable per-item block PARKS one item, it does NOT stop the track (§11.4.101 — the track backfills its next non-blocked domain item). Classification: universal (§11.4.17) — references no project hardware; the consuming project supplies its track→domain map + priority source per §11.4.35. Composes §11.4.42/.58/.72/.94/.97/.101/.103/.167/.176/.178/.179/.182/.187. Propagation gate `CM-COVENANT-114-192-PROPAGATION` (literal `11.4.192`) + recommended gate `CM-CONTINUOUS-MULTITRACK-BACKFILL` (on a track-worker completion the ruler re-assigns that track's next-highest-priority domain work-set; a track with a completed worker AND remaining actionable domain items but no new dispatch →

the gate FAILs) + paired §1.1 mutation (let a free track sit idle while its domain holds an actionable item, OR treat "merges done" as a valid idle reason → the gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.192. Non-compliance is a release blocker. No escape hatch — no `--allow-idle-track`, `--skip-backfill`, `--stop-multitrack`, `--merges-done-is-idle`, `--defer-track-reassign` flag.

§11.4.193 — Anti-blind-typing mandate: every UI interaction MUST be "seen" and "understood" via OCR/vision/screenshot proof, NEVER blind-typed (User mandate, 2026-07-13). Verbatim operator mandate: "No blind typing of any kind anywhere!!!"; "Login to apps and services MUST BE done with understanding and seeing the whole UI ALWAYS!!!"; "blind typing is STRICTLY FORBIDDEN, and instead everything MUST BE seen and understood in coordination with proper models and means we have incorporated in HelixQA!!!". Blind typing — sending keystrokes/text/input events to a UI surface WITHOUT first capturing AND understanding what is on screen (OCR / vision / uiautomator / accessibility-tree / screenshot) AND confirming the result after each interaction — is STRICTLY FORBIDDEN everywhere and always, for every agent/tool/automation/platform/context. It is LITERAL bluff: claiming "credentials entered"/"login completed" without visual proof the keystrokes landed on the intended fields/screen producing the intended result is a §11.4 PASS-bluff at the UI-interaction layer (a credential typed into the wrong field, or a WebView that silently swallowed input, = end-user feature broken while the agent reported success). (1) **SEE-BEFORE-YOU-TYPE** — before ANY `input text` / `input keyevent` / `adb shell input` / uiautomator gesture, capture the current screen state (screencap + Tesseract OCR ≥150 DPI, OR uiautomator dump grep, OR accessibility-tree dumpsys, OR HelixQA VisionEngine frame per §11.4.159/§11.4.160) AND confirm the intended target element is visible AND focused — the capture IS the §11.4.5/§11.4.69 `touch_input` evidence. (2) **VERIFY-AFTER-YOU-TYPE** — after EVERY interaction step, capture a NEW screen state + confirm the intended change (5 `input text` + 1 final screenshot = blind typing). (3) **OCR/VISION ORACLE MANDATORY** — Tesseract ≥150 DPI / uiautomator-dump grep / HelixQA VisionEngine / accessibility-tree dumpsys. (4) **WEBVIEW** — near-empty uiautomator hierarchy → OCR fallback per §11.4.117; if even OCR is infeasible (secure/DRM-blanked surface captures black) → the login is operator_attended-only, SKIP-with-reason per §11.4.3, NEVER blind-type-and-claim-success. (5)

CREDENTIAL-SAFE OCR — redact credential characters from OCR output + logs per §11.4.10; evidence = the redacted screenshot / OCR text. (6) **ANTI-BLUFF** — every login/interaction PASS cites a captured-evidence artefact containing pre-interaction state + post-interaction state + OCR/vision verification of the expected change; self-validated per §11.4.107(10) with a golden-good + golden-bad fixture pair. Classification: universal (§11.4.17). Composes §11.4 / §11.4.5 / §11.4.6 / §11.4.69 / §11.4.107 / §11.4.117 / §11.4.159 / §11.4.160 / §11.4.10 / §11.4.52 / §11.4.3. Propagation gate `CM-COVENANT-114-193-PROPAGATION` (literal `11.4.193`) + recommended gate `CM-NO-BLIND-TYPING` + paired §1.1 mutation (strip OCR verification before an `input text` → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.193. Non-compliance is a release blocker. No escape hatch — no `--blind-typing-ok`, `--skip-ocr-verification`, `--assume-input-landed`, `--login-without-proof`, `--type-without-seeing` flag.

§11.4.194 — Exhaustive all-scenario, all-angle code-review + verify-against-captured-runtime-evidence mandate (User mandate, 2026-07-14). Verbatim operator mandate: "CRITICAL: Code review MUST TAKE into the accounts all scenarios, look at the problem or change from all angles, deepest possible analisys! No gaps, shortcomings or false results or bluff is allowed! Extend our root constitution Submodule and all relevant files properly so the review process is more precise and we do not have situations like the last one!" STRENGTHENS §11.4.125/§11.4.134/§11.4.142 — every code-review MUST, before GO: (1) ENUMERATE the FULL input/scenario space of the change AND the defect (every dimension/path/interacting feature, not just the reported scenario); for a MULTI-FACTOR failure mode (a product of N terms, ANY degenerate term — zero/empty/null/default-init/out-of-range — triggers the failure) enumerate + INDEPENDENTLY VERIFY the fix covers EACH factor — verifying one and assuming the rest is the exact forbidden gap (forensic FACT "the last one": a boot-crash fix got a clean zero-finding GO that dismissed a clamp branch as an informational NIT "effectively unreachable ... always valid" WITHOUT PROVING it against the runtime input space; the target STILL crash-looped because the degenerate default-init input the review assumed impossible DID happen — a `size = A × B` product (e.g. `frame_size = channels × bytes_per_sample(format)`) with TWO independent zero-triggers, only ONE factor verified, the OTHER assumed away, the pre-fix crash evidence never cross-checked); (2) PROVE every assumption — "can't happen"/"unreachable"/"always valid"/"caller never sends X" MUST be PROVEN from code+spec+captured evidence, never on faith (§11.4.6); an unproven assumption that proves false is a §11.4 review-layer bluff of PASS-bluff severity; a NIT dismissing a path as unreachable MUST carry the proof of unreachability (a bare "effectively unreachable" NIT is itself a finding); (3) VERIFY the fix against ANY captured runtime evidence (crash log/tombstone/stack-trace/sink-probe/ §11.4.108 runtime-signature/reproducer) — confirm the fix ACTUALLY prevents the CAPTURED failure against the exact captured conditions; a GO on a fix that then fails at runtime on a scenario present in the evidence (or one the review chose not to consider) is a bluff; (4) ALL ANGLES — correctness, full input domain (every factor/boundary/degenerate value), every downstream/interacting must-not-break feature (§11.4.133 target/hardware safety + shared-state features e.g. audio/ multichannel/passthrough), the TESTS' OWN per-dimension coverage (a test checking only one factor of a multi-factor bug is itself a gap the review catches), the §1.1 mutation per dimension; (5) HONEST BOUNDARY (§11.4.6) — "exhaustive" = the enumerated + evidence-grounded scenario space, NOT literal omniscience; an un-analysable dimension is stated as an explicit gap (UNCONFIRMED/UNKNOWN/ PENDING_FORENSICS + tracked follow-up), NEVER silently assumed safe. Classification: universal (§11.4.17). Composes §11.4.125/.134/.142/.6/.108/.130/.107(10). Propagation gate `CM-COVENANT-114-194-PROPAGATION` (literal `11.4.194`) + recommended gate `CM-EXHAUSTIVE-REVIEW-ALL-SCENARIOS` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.194. Non-compliance is a release blocker. No escape hatch — no `--partial-scenario-review`, `--assume-unreachable-OK`, `--skip-captured-evidence-crosscheck`, `--single-factor-review-suffices`, `--dismiss-branch-without-proof` flag.

§11.4.195 — Branch-structure governance: taxonomy + merge-after-live-QA + flavor/product non-merge (User mandate, 2026-07-14). Verbatim operator rules: branch taxonomy `feat/<NAME_OF_FEATURE>` / `product/<NAME_OF_PRODUCT>` / `flavor/<NAME_OF_FLAVOR>`; tracks on FEATURE branches or `main` MUST — after full confirmation + validation + verification + full LIVE MANUAL TESTING (§11.4.185) — MERGE to `main` COMPLETELY (all submodules + main repo), ff-only, NEVER force-push (§11.4.113), no commit lost (§9.2); tracks with an entirely-new FLAVOR/PRODUCT do ALL work on a `flavor/` AND/OR `product/` branch (both may exist) which does NOT merge to `main` the same way; commit + push ALL branches to ALL upstreams (§2.1), ff-only. Every controlled branch belongs to one of THREE canonical taxonomy classes, each with a distinct integration lifecycle: **(A) Taxonomy** (extends §11.4.181/§11.4.167(D)) — `feat/<slug>` = FEATURE (merges to `main` after live-manual-QA, clause B); `product/<slug>` = PRODUCT line + `flavor/<slug>` = FLAVOR line (do NOT merge to `main` the standard way, clause C); `<slug>` lowercase snake/kebab (§11.4.29); the canonical branch NAME is minted ONCE + used IDENTICALLY across main repo + every touched owned submodule (§11.4.181), recorded once in the §11.4.176 claim registry / §11.4.93 `logic_group → branch` map + LOOKED UP never re-invented (§11.4.6); the prefix STRINGS are operator-canonical, a project already on a different feature-prefix (`feature/…`) is compliant so long as it is the ONE canonical name per §11.4.181, reconciled by the §11.4.181 risk-free MERGE-onto-canonical path (ff-only, no commit lost, never a rename that drops commits). **(B) FEATURE/main → merge-after-live-QA** — before merging to `main` satisfy IN ORDER automated four-layer (§11.4.4(b)) + §11.4.40 full-suite retest THEN full LIVE MANUAL TESTING by the QA team (§11.4.185 FINAL human gate, automated GREEN necessary NOT sufficient); only after §11.4.185 may the merge proceed + it MUST be COMPLETE (main repo + every touched owned submodule) ff-only onto latest `main` (§11.4.113), NEVER force-push, NO commit lost / NO history rewrite (§9.2/§11.4.41 step 3), pushed to ALL upstreams (§2.1); binds the §11.4.185 gate as an explicit precondition of the §11.4.167(I)/§11.4.188 approval-gated back-merge. **(C) FLAVOR/PRODUCT lines → do-not-merge-to-main-the-same-way** — an entirely-new flavor/product line does ALL work on a `product/<NAME>` AND/OR `flavor/<NAME>` branch (both may coexist), a divergent line with its own release lifecycle that does NOT flow back to `main` via clause B; STILL obeys §11.4.113 (no force-push) / §9.2 (no commit lost) / §2.1 (push all branches all upstreams ff-only) / §11.4.188 (MAY merge `main` INTO itself to stay close — trunk-into-line permitted; only line-into-trunk is gated/withheld). **(D) Push discipline** (§2.1/§11.4.113) — commit + push ALL branches (feature/product/flavor) to ALL upstreams ff-only, NEVER force-push; new-branch creation is inherently ff (§9.2) + non-disruptive from an existing HEAD, permitted in a dirty tree under §11.4.84 provided no commit is in-flight. Honest boundary §11.4.6 — governs branch STRUCTURE + integration DIRECTION, does NOT prove any branch's work correct (still §11.4.108 runtime-signature + §11.4.40 retest + §11.4.185 manual-QA). Classification: universal (§11.4.17) — strict generalisation of §11.4.181 + §11.4.167(D)/(I) + §11.4.188 + §11.4.185. Composes §2.1/§9/§9.2/§11.4.6/ §11.4.17/§11.4.29/§11.4.35/§11.4.40/§11.4.41/§11.4.84/§11.4.93/§11.4.113/

§11.4.167(D)/(I)/§11.4.176/§11.4.178/§11.4.181/§11.4.185/§11.4.188/§11.4.191. Propagation gate `CM-COVENANT-114-195-PROPAGATION` (literal `11.4.195`) + recommended gate `CM-BRANCH-TAXONOMY-CLASS` (every controlled branch prefix ∈ `feat/` | `product/` | `flavor/` |reserved `main` / `master` ; a `product/` | `flavor/` branch is NOT the source of a merge onto `main` ; a `feat/` | `main` back-merge cites a §11.4.185 manual-QA confirmation) + paired §1.1 mutation (a `product/*` merged onto `main` , OR a `feat/*` back-merge with no §11.4.185 citation → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.195. Non-compliance is a release blocker. No escape hatch — no `--any-branch-prefix-OK` , `--merge-product-to-main` , `--skip-live-qa-before-merge` , `--force-push-branch` , `--partial-submodule-merge` flag.

§11.4.196 — Native-alias-first priority + per-alias real-signal limit/subscription tracking + auto-rebind-on-recovery + resource-detection-by-real-identity-not-substring-match (User mandate, 2026-07-14). Verbatim operator mandate: use ALL operational claude-NATIVE aliases FIRST — as long as ≥1 native alias is operational (NOT session-rate-limited, NOT weekly-limit-reached, NOT expired-subscription) — before ANY provider alias; track per alias the session rate-limit / weekly-limit / subscription expiry+renewal from REAL signals (never faked, §11.4.6) + record when each becomes operational AGAIN; the moment a higher-priority native recovers, rebind tracks to it; and fix the host-budget guard footgun where a `pgrep -f` substring-match false-REFUSED every worker spawn. The multi-track alias-binding + rate-limit/subscription-tracking layer of the §11.4.187 ruler orchestration MUST implement ALL of: **(A) NATIVE-ALIAS-FIRST PRIORITY** — every OPERATIONAL native (`claudeN`) alias is selected before ANY provider, UNCONDITIONALLY, while ≥1 native is operational; the guarantee is STRUCTURAL — a native/provider CLASS partition scanned native-class-completely-before-provider-class (§11.4.111 resolve-by-CLASS-not-by-file-position), so a provider NEVER wins while a non-cooled native exists regardless of roster order; within a class the operator-mandated order (native equal-capability, then providers deepseek→xiaomi→opencode→kimi→…) is authoritative decision-DATA (§11.4.28) never re-invented (§11.4.6). **(B) PER-ALIAS REAL-SIGNAL LIMIT + SUBSCRIPTION TRACKING** — "operational?" per alias from REAL captured signals, NEVER faked/guessed (§11.4.6): the §11.4.187 captured 429 signature (`isApiErrorMessage` + `apiErrorStatus:429` / `api_error_status:429`) subclassified into reason-CLASS closed set `{session | weekly | subscription}` from real markers (reset/weekly-limit/session-limit); each limited alias records its CLASS + an operational-again EPOCH with a DIFFERENT window per class — session (self-heals at reset, short), weekly (~until the weekly reset, long), subscription (INDEFINITE — operational-again UNKNOWN until a REAL renewal signal, parked with a far-future sentinel; §11.4.6 NEVER guesses a renewal date, an operator/config supplies it). A limited alias is NEVER auto-selected until its epoch passes (session/weekly) or a real renewal/recovery clears it (subscription). **(C) AUTO-USE-ON-RECOVERY** — the moment a higher-priority native recovers (window elapsed / renewal marked), rebind tracks UPWARD to it (a `promote` -class op keyed off a comparable priority-rank: native class before provider, then config order), PRESERVING each track's worktree AND device leases (leases key on the STABLE track id — the alias changes, not the track; §11.4.119/§11.4.187); idempotent. **(D) RESOURCE-DETECTION BY REAL IDENTITY, NOT SUBSTRING-MATCH (RB-02 guard fix)** — a "heavy work in flight" guard MUST identify a real resource by its ACTUAL process command line read from the OS (`/proc/<pid>/cmdline`), NEVER a bare `pgrep -f REGEX` substring match that false-matches a process merely MENTIONING a token — a pattern-CARRIER (a process quoting the whole alternation pattern; a `claude` / worker whose prompt embeds the token; a launcher/monitor/test), a same-tool ENGINE process, or the guard's own process/parent (the §12.12 self/sibling pgrep footgun — forensic FACT 2026-07-13: the guard false-REFUSED EVERY §11.4.187 spawn on a host with ZERO real builds because a live `claude` worker's prompt quoted the pgrep pattern); the guard re-reads each matched PID's real cmdline + excludes those carrier/engine/worker/self classes; an UNREADABLE/ vanished cmdline is conservatively counted as the real resource (REFUSE — safe reversible default, §11.4.101/§12.8), so a genuine build is still caught. **(E) ANTI-BLUFF (§11.4/§11.4.108)** — four-layer coverage each (pre-build gate + on-host test + paired §1.1 mutation + captured evidence §11.4.69) with §11.4.115 RED-polarity + §11.4.50 determinism; honest boundary §11.4.6 — mandates the MECHANISM + REAL-signal derivation of "operational", genuine live per-subscription quota-isolation proof still needs live tokens (operator-scheduled acceptance window), never conflated with a mock/hermetic-contract GREEN. Classification: universal (§11.4.17). Composes §11.4.187/.182/.176/.111/.101/.103/.6/.108/.115/.119/§12.12/§12.6/§12.8. Propagation gate `CM-COVENANT-114-196-PROPAGATION` (literal `11.4.196`) + recommended gate `CM-ALIAS-PRIORITY-LIMIT-TRACKING` (native-first CLASS partition; per-alias reason-class + operational-again tracking; auto-rebind-on-recovery; build-guard resolves by real cmdline not bare substring) + paired §1.1 mutation (swap native/provider priority bases so a provider outranks a native → recommended gate FAILs; flip the guard's strict-cmdline-filter to the pre-fix bare `pgrep` → its RB-02 test FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.196. Non-compliance is a release blocker. No escape hatch — no `--provider-before-operational-native`, `--fake-limit-state`, `--guess-renewal-date`, `--no-rebind-on-recovery`, `--substring-match-build-guard`, `--skip-real-signal-classification` flag.

§11.4.197 — Research / kicked-off-work completion mandate: every research effort MUST be driven to full, wired, verified completion or explicitly closed — never left un-wired in the backlog (User mandate, 2026-07-15). Verbatim operator mandate: "all research work we start MUST BE fully finished, completely confirmed, validated and verified! It MUST NOT happen for it to stay at the backlog uncompleted and un-wired! Loss of requirements, failure to incorporate them is FORBIDDEN!!! Everything MUST BE fully completed! No exceptions!" Every research effort / kicked-off improvement / design doc / spike / prototype / PoC / partially-landed feature a project STARTS MUST be driven to FULL completion OR explicitly evidence-backed CLOSED — there is no third "sitting un-wired in the backlog" state; "we researched/designed/started it" is NEVER terminal, and a `docs/research/*` doc / design doc / spike / half-wired feature that is neither COMPLETED-and-wired nor explicitly-closed is a §11.4.197 violation of §11.4 PASS-bluff severity at the requirements-integrity layer (the requirement was accepted, work began, then it silently evaporated); loss of requirements + failure to incorporate started work are FORBIDDEN. **COMPLETED** = (1) implemented + (2) WIRED (integrated AND used by default, NOT merely present behind a dead flag / uncalled function / unreferenced module / never-selected path — mere-presence-without-wiring is the §11.4.124 dead-code + §11.4.108 layer-2/3 SOURCE→ARTIFACT→RUNTIME gap) + (3) confirmed/validated/verified with captured physical evidence §11.4.5/§11.4.69/§11.4.107 on a clean deployment §11.4.108 runtime-signature (never metadata-only/config-only/absence-of-error/grep-without-runtime §11.4/§11.4.1) + (4) tracked to a terminal Status §11.4.15/§11.4.33. **Explicitly CLOSED** = deliberately terminated with a documented evidence-backed reason from the closed vocabulary — §11.4.90 Obsolete / §11.4.112 structurally-impossible / §11.4.122 operator-confirmed drop — NEVER silent abandonment ("no time"/"lower priority" = §11.4.21 Operator-blocked or a still-open tracked item, NOT a close). **Mechanical tracker binding (§11.4.93):** every research effort / design doc / kicked-off improvement MUST map to ≥1 tracked workable item (ATM-NNN §11.4.54) whose completion the loop enforces; a research doc with NO implementing item, or a tracked item stalled un-wired past the cadence, is a violation; the autonomous loop (§11.4.87/.94/.97/.103/.126/.192) MUST treat every un-wired research effort as an actionable queue item — driving it to COMPLETED or explicitly CLOSED is exactly the "keep working until the queue is genuinely empty" contract. Honest boundary §11.4.6 — "fully completed" is bounded by wired+verified OR evidence-backed-closed, NOT "every speculative idea ships"; every STARTED effort reaches a documented terminal state, proven never assumed. Classification: universal (§11.4.17). Composes §11.4.8/.87/.94/.97/.103/.118/.123/.124/.150/.192/.90/.112/.122/.21/.93/.54/.108/.15/.33/.126. Propagation gate `CM-COVENANT-114-197-PROPAGATION` (literal `11.4.197`) + recommended gate `CM-RESEARCH-COMPLETION-TRACKED` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.197. Non-compliance is a release blocker. No escape hatch — no `--research-may-stay-unwired`, `--skip-completion-tracking`, `--abandon-without-close`, `--present-but-unwired-is-done`, `--backlog-research-OK`, `--lose-requirement-OK` flag.

§11.4.198 — Default working mechanisms: multi-alias native-first orchestration AND heavy token-optimization are ALWAYS-ON defaults for single- and multi-track work (User mandate, 2026-07-15). Operator mandate: multiple claude-toolkit aliases (native-first) AND heavy token-optimization MUST be the DEFAULT mechanisms, always used with single- and multi-track development and with single or multiple working agents. Two proven mechanisms promoted to STANDING always-on defaults, engaged automatically for EVERY configuration (single- OR multiple-agent, single- OR multi-track), never a per-request opt-in: **(A) MULTI-ALIAS NATIVE-FIRST ORCHESTRATION** — the claude-toolkit multi-alias mechanism (every OPERATIONAL claude-native alias before ANY provider per §11.4.196 native-alias-first priority + per-alias real-signal limit/subscription tracking + auto-rebind-on-recovery, driven by the §11.4.187 automatic multi-track ruler) MUST be the default execution substrate for ALL work, NOT only the multi-track case; a SINGLE-track/SINGLE-agent session STILL binds its worker through the same native-first alias-priority mechanism (best operational native + rebind-on-recovery), a MULTI-track/multiple-agent config uses it across every track (§11.4.176

exactly-once claim + §11.4.182 track+branch+alias label); running on an arbitrary/hardcoded/provider-first alias when an operational native exists, or bypassing the mechanism because "it is just one agent", is a violation. **(B) HEAVY TOKEN-OPTIMIZATION** — the §11.4.141 token-efficiency discipline (compact-anchor layout with full-one-hop-away, targeted reads over whole-file dumps, subagent context isolation §11.4.20/§11.4.70, evidence-hash references not evidence dumps) MUST apply ALWAYS — every read/dispatch/doc-edit — for single- AND multi-track, single OR multiple agents; token-heavy patterns (re-reading whole large files when a targeted slice suffices, dumping full evidence into context, monolithic non-subagent execution of parallelisable work) when a token-efficient path exists are violations. Honest boundary §11.4.6 — "default" = auto-engaged without an operator request; it does NOT weaken host-safety, the multi-alias/multi-agent fan-out stays under the §12.6 60% memory ceiling + §12.12 thread-headroom + §11.4.58 parallel-agent cap (contending/redundant agents are NOT "the default", §11.4.183 maximum-USEFUL-throughput); §11.4.198 BINDS the two mechanisms as always-on, it does NOT redefine them (mechanics live in §11.4.196/§11.4.187/§11.4.141). Classification: universal (§11.4.17). Composes §11.4.141/.187/.196/.182/.176/.183/.103/.58/.20/.70/.126/§12.6/§12.12. Propagation gate `CM-COVENANT-114-198-PROPAGATION` (literal `11.4.198`) + recommended gate `CM-DEFAULT-MECHANISMS-ALWAYS-ON` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.198. Non-compliance is a release blocker. No escape hatch — no `--single-agent-skips-alias-priority`, `--provider-first-default`, `--token-optimization-optional`, `--heavy-tokens-OK`, `--not-default-mechanism` flag.

§11.4.199 — Exact-reproduction-sequence mandate: when a working reproduction exists, the investigation MUST use ITS exact sequence — a deviating repro proves NOTHING (research-derived, 2026-07-15). Forensic FACT (2026-07-15): an investigation concluded a defensive "cover" mechanism "never engages" (and the defect it guards was therefore unfixed) — the conclusion was WRONG because the hand-rolled reproduction drove the target with a DIFFERENT input than the existing test (a HOME key where the test sent a media-pause key), so it never reached the PRECONDITION the mechanism keys off; re-driven with the TEST'S EXACT sequence the mechanism provably ENGAGED. Every investigation / debugging pass / validation that reproduces a defect or exercises a code path MUST: (1) **use the existing reproduction's EXACT sequence** when a working reproduction already exists (a test/harness/recorded sequence that demonstrably reaches the defect) — same inputs, same order, same timing, same preconditions, same topology (§11.4.3); a hand-rolled approximation is NOT a substitute; (2) **a deviating variant MUST FIRST prove it reaches the precondition** with captured evidence before ANY conclusion is drawn from it; (3) **a repro that misses the precondition proves NOTHING** — it is NOT evidence of absence (§11.4.6), and concluding "defect present"/"defect absent"/"the mechanism never engages" from it is a §11.4/§11.4.1 bluff at the INVESTIGATION layer (a FAIL-bluff when it wrongly condemns working code, a PASS-bluff when it wrongly clears broken code); (4) **record the exact driving sequence as captured evidence** so it is replayable + auditable, every deviation explicit + justified + precondition-proven. STRENGTHENS §11.4.102 (Phase-2 reproduce) / §11.4.146 (reproduce-first) / §11.4.115 (RED on the broken artifact) — they mandate THAT you reproduce, §11.4.199 mandates HOW: with the sequence PROVEN to reach the defect. Classification: universal (§11.4.17). Composes §11.4.1/.3/.5/.6/.7/.102/.107/.115/.123/.135/.146. Propagation gate `CM-COVENANT-114-199-PROPAGATION` (literal `11.4.199`) + recommended gate `CM-EXACT-REPRO-SEQUENCE` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.199. Non-compliance is a release blocker. No escape hatch — no `--approximate-repro-OK`, `--skip-precondition-proof`, `--hand-rolled-repro-suffices`, `--absence-proves-absence`, `--deviating-sequence-conclusion-OK` flag.

§11.4.200 — Non-targetable deploy/flash tooling MUST isolate exactly ONE eligible target and VERIFY-AFTER-WRITE on the INTENDED target (research-derived, 2026-07-15). Forensic FACT (2026-07-15): a flashing tool with no device-selection argument AUTO-SELECTED a device from the bus, silently wrote the freshly-built image to the WRONG board (a sibling still attached), printed SUCCESS, and the INTENDED board kept its OLD build — every downstream test then validated a target that had NEVER received the artifact (the §11.4.108 ARTIFACT→RUNTIME layer silently broken while every tool reported green). Whenever a deploy / flash / install / write tool CANNOT address a specific target, the procedure MUST: (1) **targetability is a PRECONDITION** — EITHER pass an explicit STABLE selector (serial/UUID/stable id per §11.4.111, never an enumeration ordinal) OR GUARANTEE exactly ONE eligible target is present at the moment of the write (siblings disconnected/powered-down/excluded), the enumerated eligible-target set captured as evidence (size == 1); (2) **the tool's own success message is NOT proof** — an "ok"/exit-0 proves only that bytes reached the device the TOOL selected, NEVER the device the OPERATOR intended (§11.4.6); (3) **VERIFY-AFTER-WRITE on the intended target** — read back the deployed artifact's IDENTITY (build id / version / checksum / fingerprint) FROM the intended target and assert it equals the artifact just written; a mismatch or the OLD identity is a FAIL, never a warning (the §11.4.108 ARTIFACT→RUNTIME assertion applied to the deployment step); (4) **no validation on an unverified target** — any test run against a target whose deployed-artifact identity was NOT verified is INVALID and any PASS is a §11.4 PASS-bluff. Honest boundary §11.4.6 — isolation + read-back prove the INTENDED target holds the INTENDED artifact, NOT that the artifact is correct (§11.4.108/.40/.130). Classification: universal (§11.4.17). Composes §11.4.6/.46/.108/.111/.119/.130/.133/.139. Propagation gate `CM-COVENANT-114-200-PROPAGATION` (literal `11.4.200`) + recommended gate `CM-DEPLOY-TARGET-ISOLATED-AND-VERIFIED` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.200. Non-compliance is a release blocker. No escape hatch — no `--auto-select-target-OK`, `--trust-tool-success-message`, `--skip-post-write-verify`, `--validate-unverified-target`, `--multiple-targets-on-bus-OK` flag.

§11.4.201 — Every guard/gate MUST assert the REAL condition: a false-positive refusal is a FAIL-bluff, a false-negative pass is a PASS-bluff (research-derived, 2026-07-15). Forensic FACT (2026-07-15, two independent instances in one cycle): (a) a host-budget guard used a bare `pgrep -f '<pattern>'` and REFUSED every worker spawn on a host with ZERO real builds because a live worker's own command line merely QUOTED the pattern (the §11.4.196(D)/§12.12 carrier footgun); (b) a host-safety pre-flight refused a build on a "swap full" reading that did not reflect a real memory-pressure condition, clearing only after the operator manually cycled swap — BOTH guards refused on a PROXY signal that was NOT the real condition. Every GUARD / GATE / PRE-FLIGHT (host-safety, resource-budget, readiness, lock, capability, availability) MUST assert the REAL condition it CLAIMS, resolved from the AUTHORITATIVE source (the real `/proc/<pid>/cmdline`, the real cgroup/meminfo pressure metric, the real lock-holder liveness via `kill -0`, the real device identity), NEVER from a proxy signal that something which is NOT the condition can satisfy: (1) **a FALSE-POSITIVE REFUSAL is a FAIL-bluff (§11.4.1)** — blocking work when the condition is ABSENT is exactly as forbidden as passing while a defect is present (it halts real work AND teaches operators/agents to bypass guards); (2) **a FALSE-NEGATIVE PASS is a PASS-bluff (§11.4)**; (3) **SELF-VALIDATED (§11.4.107(10))** — every guard ships a golden-TRUE fixture (condition present → guard FIRES) AND a golden-FALSE fixture INCLUDING the known CARRIER/decoy case (condition absent but proxy signal present → guard MUST NOT fire), both wired into the meta-test; a guard firing on its golden-FALSE fixture is ITSELF the defect; (4) **conservative-safe default on an unresolvable signal (§11.4.101/§12)** — refuse the destructive/high-blast-radius action and say so HONESTLY ("could not resolve the condition" is NOT a false positive; claiming a condition never verified IS one, §11.4.6); (5) **every refusal reports its RESOLVED evidence** (real pid+cmdline, real metric+threshold, real lock holder) so a false positive is diagnosable in one step. GENERALISES §11.4.196(D) (resource-detection-by-real-identity — to EVERY guard) + §11.4.180 (liveness PROVEN before reaping — to every guard's condition). Classification: universal (§11.4.17). Composes §11.4.1/.6/.101/.107(10)/.109/.111/.180/.196(D)/§12/§12.6/§12.8/§12.12. Propagation gate `CM-COVENANT-114-201-PROPAGATION` (literal `11.4.201`) + recommended gate `CM-GUARD-ASSERTS-REAL-CONDITION` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.201. Non-compliance is a release blocker. No escape hatch — no `--proxy-signal-OK`, `--substring-match-guard`, `--false-positive-refusal-acceptable`, `--unvalidated-guard`, `--refuse-without-evidence` flag.

EXTENSIONS landed 2026-07-15 (existing anchors strengthened — full text in `Constitution.md`). §11.4.89(G) subagent long-op discipline + conductor-owns-inherently-long-ops: a SUBAGENT MUST NOT run a single synchronous op exceeding the runtime's no-progress stream watchdog (~5 min) — it backgrounds + polls (each poll IS the progress the watchdog observes); inherently-long ops (full builds / full suites / device flashes) are the CONDUCTOR's to run (no such watchdog); a stall-killed subagent is a CRASH per §11.4.147 (respawned, never absorbed as a completion); watchdog budget read from the runtime, never guessed (§11.4.6). **§11.4.106(F)** write-seam hook enforcement: the doc/DB sync MUST be enforced at the COMMIT seam (hook refuses a commit whose staged set touches a chain source while its exports/DB are stale) + the BUILD seam (pre-build gate) + the CONSTITUTION-PULL seam (§11.4.164 post-update hook) — a stale CONTINUATION (§12.10) / session-resumption file (§11.4.131) / memory artefact while live state moved is a violation; honest boundary §11.4.6 — seam hooks give eventual-consistency-at-every-write, NOT literal real-time (that needs a watch daemon, carried as a tracked §11.4.197 upgrade path, never claimed as shipped). **§11.4.147(e)** API-QUOTA / RATE-LIMIT EXHAUSTION is a FIRST-CLASS CRASH

CLASS: an agent killed mid-work by a weekly/session/subscription quota boundary is `crashed` (never `complete`) — record the REAL limit signal + reason-class per §11.4.196(B), REBIND to the next OPERATIONAL alias per §11.4.196(A)/(C) and RESPAWN there (bounded-PARK per §11.4.101 only when EVERY alias is cooled), RESUME from the killed agent's captured transcript + partial state at its EXACT last point, and the §11.4.87/.94/.97/.126 loop done-condition MUST NOT read satisfied while any quota-killed entry is non-`complete`. **§11.4.196(F)** CONFIGURED ≠ IN USE: forensic FACT — the alias orchestrator was installed with a valid config yet its binding registry was EMPTY while THREE live workers ran outside it, so `promote` rebound ZERO tracks; EVERY spawn path MUST bind THROUGH the orchestrator, and the mechanism's readiness is its RUNTIME SIGNATURE (§11.4.108) — **the binding-registry rows MUST EQUAL the live-worker set** — an empty/partial registry alongside live workers is a FALSE-READY state and claiming "ready/in use" on config alone is a §11.4/§11.4.6 orchestration-layer bluff.

§11.4.142 + §11.4.194 re-affirmed (2026-07-15 audit, no change required): the operator re-insisted 3× this session that EVERY change passes an exhaustive code review, ALWAYS, before build/commit. Audited: §11.4.142 already binds EVERY change to ANY governed repository — "no exception ... no 'just a doc edit' exemption, no 'just a one-liner' exemption, no §11.4.92-self-review substitution, no 'trivial change' carve-out ... BEFORE it is accepted, committed, or built" — with independence load-bearing (§11.4.70/§11.4.20); §11.4.194 supplies the exhaustive all-scenario / prove-every-assumption / cross-check-captured-runtime-evidence precision bar; §11.4.134 iterates the review to a ZERO-finding, ZERO-warning GO; §11.4.125 places the gate before the pre-build sweep + main build. NO carve-out exists in any of them — the family is complete as written and needed no widening this round (§11.4.6: audited, not assumed).

§11.4.202 — Reporting directives: a report MUST auto-create a fully-populated, fully-synced workable item — never a prose acknowledgement (User mandate, 2026-07-15). Verbatim operator mandate: introduce reporting directives `ISSUE` / `BUG` / `TASK` that create a proper workable item filled with all details, with the WHOLE workable-items system automatically up to date and fully in sync — the main DB, all documents, and all external tracking systems; universal, reusable on every project, fully decoupled, recognized immediately by every project that pulls the constitution. Every governed project MUST expose REPORTING DIRECTIVES — the §11.4.140 registered actions **`BUG`** (Type=Bug), **`TASK`** (Type=Task), and **`ISSUE`** (generic report, CLASSIFIED into the §11.4.16 closed set) — such that a plain-language report from the operator is TURNED INTO a real, tracked, fully-synced workable item, automatically, every time. A report that is discussed, acknowledged, answered in prose, or "noted" WITHOUT landing a tracked item is a §11.4.202 violation of §11.4 PASS-bluff severity at the requirements-intake layer: the requirement was accepted and then silently evaporated (the §11.4.197 loss-of-requirements failure, at its entry point). **(1) GRAMMAR (§11.4.140).** The three directives are ordinary registry rows, so EVERY §11.4.140 form works out of the box (`BUG :: <text>`, `DEFAULT::BUG :: <text>`, `/DEFAULT::BUG <text>`, `BUG ---> <text>`), PLUS a SIXTH form this anchor adds — the **single-colon form** **`NAME: <text>`** (`ISSUE: subtitles are late`), the shape operators actually type. The single-colon form is **REGISTERED-ACTION-ONLY by construction**: a single-colon token that is NOT a registered action is a NO-OP (an ordinary prompt), NEVER an ASK and NEVER a sub-system shortcut — because ordinary English prose (`TODO:`, `WARNING:`, `FIXME:`) shares that shape (`NOTE:` is now a registered §11.4.140 severity marker), and routing it to the §11.4.66 clarify path would question every such sentence (§11.4.6 honest boundary: the other five forms keep their ASK-on-unknown behaviour — their shapes do not occur in prose). Host-command collisions are handled by the EXISTING §11.4.140 conflict mechanism (`slash_bare: auto` + `slash_conflicts: [..]`), which already satisfies the operator's explicit-prefix requirement: a bare `/NAME` that collides with a built-in (e.g. the documented Claude Code built-in `/bug`) is NOT honored, while the namespaced `/DEFAULT::NAME` form is ALWAYS unambiguous — and the `NAME:` / `NAME ::` / `NAME --->` forms are never affected by any host command. **(2) THE ITEM.** The created item MUST carry, on ALL surfaces: a valid Status (§11.4.15, `Queued` at intake) + a Type from the CLOSED set {Bug | Feature | Task} (§11.4.16 — `ISSUE` is a REPORTING CHANNEL, never a fourth type; inventing one is a violation) + a stable auto-incremented id (§11.4.54), and a COMPREHENSIVE structured description (§11.4.148 D2 / §11.4.171): what / affected-scope / reproduction / acceptance. An undetermined section is recorded as an explicit `UNKNOWN:` gap — NEVER invented (§11.4.6). For `ISSUE`, the type is classified from the report's content and stated as FACT; if the content does not determine it, the operator is ASKED (§11.4.66 / §11.4.105) BEFORE creation; ONLY when running autonomously where asking is impossible (§11.4.101) may the §11.4.16 lowest-stakes ambiguity default `Task` be used — and then the defaulted classification MUST be recorded verbatim in the item and surfaced for reclassification, never silently asserted. **(3) THE FULL SYNC (the load-bearing half).** Creating the item is NOT sufficient: the directive MUST drive the WHOLE workable-items system into sync, in one shot — (a) the item lands in the SQLite single-source-of-truth (§11.4.93 / §11.4.95); (b) EVERY derived document is regenerated FROM the DB — the open + closed trackers, their summaries, and their `.md` / `.html` / `.pdf` / `.docx` siblings (§11.4.12 / §11.4.53 / §11.4.65 / §11.4.106); (c) the item is pushed to EVERY configured external tracker (§11.4.148 D5). **(4) NEVER A FAKED PUSH (§11.4.10 / §11.4.6 / §11.4.3).** A tracker whose credentials or whose client are absent is SKIPPED with an honest machine-readable reason (`credentials_absent` — names of the unset variables ONLY, never a value; `tracker_client_absent` — PENDING-OPERATOR-INPUT). A tracker PASS may be recorded ONLY after a real push command really exited 0. Fabricating, assuming, or silently omitting a push is a §11.4 bluff at the integration layer; a missing tracker NEVER blocks the item from landing in the DB + docs (the report is never lost because a mirror is unreachable). **(5) DECOUPLED ENGINE, CONSUMER-OWNED DATA (§11.4.28 / §11.4.177).** The item-creation + full-sync engine lives in the constitution submodule (`constitution/scripts/reporting/report_item.sh`), is inherited BY REFERENCE (never copied), and carries ZERO project literals; every project-specific value — DB path, id prefix, sync command, tracker commands + their required env vars — comes from a CONSUMER-OWNED config file (DATA, never an engine edit). The engine FAILS CLOSED with an actionable message when that config is absent — it never guesses a project's paths (§11.4.6). **(6) ANTI-BLUFF (§11.4 / §11.4.107(10)).** The mechanism ships four-layer coverage: a pre-build gate (`CM-REPORTING-DIRECTIVES`) whose FUNCTIONAL invariant SOURCES and RUNS the parser (never a grep-only assertion, §11.4.108), a paired §1.1 mutation test that proves each invariant is genuinely load-bearing, and an end-to-end suite run against a REAL SQLite DB with REAL binaries (no mocks, §11.4.27) that asserts a real row lands with the right Type/Status/id, that the docs are regenerated FROM the DB, and that an absent tracker SKIPs honestly — self-validated with a golden-good + golden-bad + negative-control fixture set (an oracle that passes its golden-bad fixture is itself the bluff). Classification: universal (§11.4.17) — the consuming project supplies its DB path, id prefix, sync command, and tracker bindings per §11.4.35. Composes §11.4.140 (grammar) / §11.4.93 / §11.4.95 (DB SSoT) / §11.4.15 / §11.4.16 / §11.4.54 (status + type + id) / §11.4.148 / §11.4.171 (comprehensive description) / §11.4.106 / §11.4.12 / §11.4.53 / §11.4.65 (doc sync) / §11.4.10 (credentials) / §11.4.66 / §11.4.105 (clarify, never guess) / §11.4.101 (autonomous decision) /

§11.4.197 (a started requirement never evaporates) / §11.4.28 / §11.4.177 (decoupling) / §11.4.164 (auto-registration on constitution pull) / §11.4.107(10) / §1.1. Propagation gate `CM-COVENANT-114-202-PROPAGATION` (literal `11.4.202`) + recommended gate `CM-REPORTING-DIRECTIVES` (registry declares ISSUE/BUG/TASK + the single-colon form is registered-only; the engine exists, is parse-clean, and is project-literal-free; the three directives really expand at runtime while prose does not; the honest-skip reasons exist and a tracker PASS is gated on a real exit 0) + paired §1.1 mutation (strip an action row, flip the single-colon form to non-registered-only, strip the parser branch, remove the prose NO-OP branch, make the engine fake a tracker push, or couple the engine to one project → the gate FAILs; strip the literal → the propagation gate FAILs).

Canonical authority: constitution submodule `Constitution.md` §11.4.202.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--report-without-item`, `--prose-acknowledgement-OK`, `--skip-item-creation`, `--skip-tracker-sync`, `--fake-tracker-push`, `--invent-a-fourth-type`, `--hardcode-project-in-engine` flag.

§11.4.207 — Instant multi-stream resume engine: whole-fleet continuation state is a durable, content-addressed, atomically-committed snapshot resumed in O(changed) reads, not an O(history) re-read (User mandate, 2026-07-15). Compact summary: the continuation mechanism (§12.10 / §11.4.127 / §11.4.131 / §11.4.205) is EXTENDED — never replaced — by a mechanical engine so a fresh session resumes ALL work (every Track, main stream, and agent) at once, INSTANTLY and at token cost O(changed streams), by reading ONE snapshot manifest + one compact blob per stream, NEVER re-reading the full handoff history. (1) **Content-addressed Merkle store** — each stream's canonical state bytes are named by their sha256 (dedup: unchanged streams share a hash → a global fleet snapshot costs O(changed)); a snapshot is ONE manifest referencing every stream's content-hash (§11.4.205(4)). (2) **O(streams) resume, O(history) NEVER on the resume path** — `RestoreAll` reads the manifest + one blob per stream; the append-only event ledger is audit + lost-update recovery only, never read to resume (§11.4.127). (3) **Tamper-evident + deterministic** — every blob is integrity-checked on read (bytes MUST hash to their id) and a deterministic round-trip (re-serialize must byte-match), no wall-clock in the hashed content, verified by a SELF-VALIDATING oracle: golden-good PASS + golden-bad detected FAIL + negative-control (a legitimately-older-but-valid snapshot) PASS — the false-positive guard proving the verifier distinguishes a LAGGING copy from a TAMPERED one (§11.4.201/§11.4.107(10)/§11.4.206(3)). (4) **Durable + atomic + single-writer + crash-safe** — atomic ref write (temp→fsync→rename→dir-fsync, §11.4.205(6)), append-only ledger (§11.4.205(6)), single-writer-per-stream refusal (§11.4.206), advisory lock with provably-stale reap (`kill -0` liveness, never steals a live lock, §11.4.180/§9.2). (5) **Project-agnostic engine, consumer-owned data** — the engine (`github.com/vasic-digital/continuum` , GitLab mirror) carries ZERO project literals, fails closed rather than guess a store path (§11.4.6/§11.4.177), and is inherited BY REFERENCE as a depth-1 submodule under the §11.4.28(C) carve-out (`constitution/submodules/continuum/` , `helix-deps.yaml` , zero own-org deps); every consumer supplies its store path / actor / stream naming as DATA. Anti-bluff (§11.4/§11.4.108): the whole engine ships four-layer coverage (unit + integration + e2e + stress + chaos §11.4.85, race-clean) with a measured resume Metric (`blobs_read == 1 + streams` is the O(streams) proof) and the self-validating oracle as captured evidence (§11.4.5/§11.4.69/§11.4.107); honest boundary (§11.4.6): it snapshots the state an agent REPORTS, not its execution image (not a CRIU/durable-execution runtime, §11.4.112), and `est_tokens` is a documented `bytes/4` heuristic, not a tokenizer count. Classification: universal (§11.4.17). Composes §12.10 / §11.4.127 / §11.4.131 / §11.4.205 / §11.4.116 / §11.4.147 / §11.4.180 / §11.4.201 / §11.4.206 / §11.4.107(10) / §11.4.28 / §11.4.177 / §11.4.85 / §9.2. Propagation gate `CM-COVENANT-114-207-PROPAGATION` (literal `11.4.207`) + recommended gate `CM-CONTINUUM-RESUME-ENGINE-PRESENT` (the engine present under `constitution/submodules/continuum/` , `go test -race ./...` green, `continuum selfcheck` = good=PASS/bad=FAIL/negctrl=PASS, zero project literals) + paired §1.1 mutation (strip the literal → propagation gate FAILs; downgrade the resume path to read the ledger, or the oracle to pass its golden-bad → the mechanism gate FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.207. Non-compliance is a release blocker. No escape hatch — no `--resume-reads-history` , `--skip-integrity-check` , `--unvalidated-oracle-OK` , `--hand-typed-resume-state` , `--per-project-resume-reimpl` flag. **§11.4.208 — Operator-request-history document: every project maintains a project-local, always-in-sync ledger of every operator request/prompt (User mandate, 2026-07-15).** Compact summary: every governed project MUST maintain a single PROJECT-LOCAL, append-only, newest-first request-history document (e.g. `docs/requests/history.md` , path declared once per §11.4.35) recording EVERY operator request/prompt with EXACTLY these operator-mandated fields — (1) full request content (verbatim where archived, else a faithful complete summary, marked which); (2) accepted timestamp with an EXPLICIT timezone stamp on every entry (project-declared default zone, `time UNKNOWN` where only a date is recoverable, never fabricated); (3) the Track that processed it (§11.4.176/.178/.182); (4) the live alias that took the item (§11.4.182/.196); (5) the model + effort used. (A) NO INVENTION (§11.4.6) — an unrecoverable field is recorded literally `UNKNOWN`, never guessed (a fabricated timestamp/alias/model is a §11.4/§11.4.1 requirements-record bluff); (B) honest reconstruction boundary for pre-mandate sessions (best-effort from durable dated records, unrecoverable fields `UNKNOWN`, boundary stated explicitly); (C) always in sync (§12.10) + §11.4.44 revision header + four-format export §11.4.65/§11.4.73; (D) keep-applying mechanism (a helper script and/or a `UserPromptSubmit`-class hook appends a newest-first row per new prompt with deterministic track/alias derivation §11.4.182, honest `?` / `UNKNOWN` when undeterminable; a helper-without-hook is honestly partial, hook-wiring a tracked §11.4.197/§12.10 follow-up, never claimed as automatic capture it does not perform); (E) complements NEVER replaces the per-session no-loss ledger (§11.4.197/.202), the workable-items DB SSoT (§11.4.93/.95), and the §12.10/§11.4.131 session-resumption artefacts (loss of a requirement at intake FORBIDDEN §11.4.197); (F) RULE-UNIVERSAL / DOCUMENT-PROJECT-LOCAL (§11.4.35) — this universal RULE lives in the constitution submodule, the DOCUMENT + its path + the default timezone are consumer-owned DATA and are NEVER placed into the constitution (§11.4.28). Classification: universal (§11.4.17). Composes §11.4.6/.35/.44/.65/.73/.118/.131/§12.10/.176/.178/.182/.196/.197/.202/.93/.95. Propagation gate `CM-COVENANT-114-208-PROPAGATION` (literal `11.4.208`) + recommended gate `CM-REQUEST-HISTORY-DOC-PRESENT` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.208. Non-compliance is a release blocker. No escape hatch — no `--request-history-optional`, `--skip-request-log`, `--no-timezone-stamp`, `--invent-missing-field`, `--request-history-in-constitution`, `--prompt-without-history-entry` flag.

§11.4.209 — Code-review MUST run on the Fable model at xhigh effort (Opus xhigh fallback) (User mandate, 2026-07-15). Compact summary: every mandatory independent code-review the constitution requires — the §11.4.125 code-review-agent gate before pre-build + main build, the §11.4.142 universal every-change-review, and the §11.4.194 exhaustive all-scenario/all-angle review, iterated to a zero-finding GO per §11.4.134 — MUST be dispatched on the **Fable model at `xhigh` effort**; (A) fallback (closed, ordered) — ONLY when Fable is genuinely unavailable (FACT per §11.4.6, never guessed) use **Opus at `xhigh` effort**; a lower-effort or a non-Fable/non-Opus model is NOT an acceptable review substrate, `xhigh` REQUIRED in both cases; (B) independence preserved (§11.4.70/.20) — §11.4.209 sets the MODEL+EFFORT, not the independence / iterate-until-GO §11.4.134 / captured-evidence §11.4.5/.69/.107; (C) honest boundary §11.4.6 — a review claimed done that ran on a weaker model/lower effort than Fable-xhigh (or Opus-xhigh on genuine Fable-unavailability) is a §11.4 review-substrate bluff; the review's captured evidence records which model+effort actually performed it (and, on fallback, the Fable-unavailability fact). STRENGTHENS §11.4.125/.142/.194/.134, does not replace them. Classification: universal (§11.4.17). Composes §11.4.125/.134/.142/.194/.6/.20/.70/.196. Propagation gate `CM-COVENANT-114-209-PROPAGATION` (literal `11.4.209`) + recommended gate `CM-CODE-REVIEW-FABLE-XHIGH` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.209. Non-compliance is a release blocker. No escape hatch — no `--review-any-model`, `--review-lower-effort-OK`, `--skip-fable`, `--non-xhigh-review`, `--fallback-without-fable-unavailability-proof` flag.

§11.4.208 — Operator-request-history document: every project maintains a project-local, always-in-sync ledger of every operator request/prompt (User mandate, 2026-07-15). Verbatim operator mandate: introduce a request-history document capturing per operator request/prompt — full content, accepted-timestamp with explicit timezone, the track that processed it, the alias that took the workable item, and the model + effort used — always in sync, no false data, no bluff; the operator's load-bearing correction (§11.4.35/§11.4.17): the RULE is universal (→ constitution submodule), the request-history DOCUMENT itself is PROJECT-LOCAL (never placed into the constitution). Every governed project MUST maintain a single project-local, append-only, newest-first **operator-request-history document** at a fixed project-declared path (§11.4.35, never moved without a §11.4.66 operator decision), each entry carrying EXACTLY: (1) **Request content** (verbatim where archived, else a faithful complete summary, marked which); (2) **Accepted (when)** — date + time + TIMEZONE stated EXPLICITLY on every entry (consuming project declares its default zone §11.4.35; date-only ⇒ `time UNKNOWN`, never fabricated); (3) **Track** (§11.4.176/.178/.182); (4) **Alias** (§11.4.182/.196); (5) **Model + effort**. (A) **NO INVENTION (§11.4.6)** — an unrecoverable field is recorded literally `UNKNOWN`, never guessed; a fabricated timestamp/alias/model is a §11.4/§11.4.1 bluff at the requirements-record layer. (B) **HONEST RECONSTRUCTION BOUNDARY** — pre-mandate sessions are best-effort reconstructed from the durable dated record with unrecoverable fields marked `UNKNOWN` and the boundary stated explicitly, never silently implied complete. (C) **ALWAYS IN SYNC + EXPORTED** — §11.4.44 revision header, §12.10 append-on-every-new-prompt, §11.4.65/§11.4.73 four-format export (`.md` + `.html` + `.pdf` + `.docx` where applicable, siblings ≥ `.md`). (D) **KEEP-APPLYING MECHANISM** — a helper script and/or a `UserPromptSubmit`-class hook appends a newest-first row per NEW prompt at capture time (track/alias derived deterministically §11.4.182, honest `?` / `UNKNOWN`); a helper landed WITHOUT the auto-capture hook is honestly a partial mechanism (hook wiring = tracked §11.4.197/§12.10 follow-up, never claimed as auto-capture it does not perform). (E) **COMPLEMENTS, NEVER REPLACES** any per-session no-loss ledger (§11.4.197/.202), the workable-items DB SSoT (§11.4.93/.95), or the §12.10/§11.4.131 resumption artefacts; requirement loss at intake is FORBIDDEN (§11.4.197). (F) **RULE-UNIVERSAL / DOCUMENT-PROJECT-LOCAL (§11.4.35)** — the RULE lives in the constitution submodule; the document + its path + default timezone are consumer-owned project DATA (§11.4.28 decoupling), NEVER placed into the constitution. Classification: universal (§11.4.17). Composes §11.4.6/.35/.44/.65/.73/.118/.131/.176/.178/.182/.196/.197/.202/.93/.95/§12.10. Propagation gate `CM-COVENANT-114-208-PROPAGATION` (literal `11.4.208`) + recommended gate `CM-REQUEST-HISTORY-DOC-PRESENT` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.208. Non-compliance is a release blocker. No escape hatch — no `--request-history-optional`, `--skip-request-log`, `--no-timezone-stamp`, `--invent-missing-field`, `--request-history-in-constitution`, `--prompt-without-history-entry` flag.

§11.4.209 — Code-review MUST run on the Fable model at xhigh effort (Opus xhigh fallback) (User mandate, 2026-07-15). Verbatim operator mandate: the mandatory independent code-review MUST ALWAYS be performed with the **Fable model at xhigh effort**; if Fable is unavailable, use **Opus at xhigh effort**. Every mandatory independent code-review the constitution requires — the §11.4.125 code-review-agent gate before pre-build + main build, the §11.4.142 universal every-change-review (no exception), the §11.4.194 exhaustive all-scenario/all-angle review, iterated to a zero-finding GO per §11.4.134 — MUST be dispatched on the

Fable model at `xhigh` effort. (A) **Fallback (closed, ordered)** — ONLY when Fable is genuinely unavailable (established as FACT §11.4.6, never a guessed unavailability) the review runs on **Opus at `xhigh` effort**; a lower-effort or a non-Fable/non-Opus model is NOT an acceptable review substrate, `xhigh` is REQUIRED in both cases (the §11.4.194 all-angle depth is unreachable lower). (B) **Independence preserved (§11.4.70/.20)** — the Fable/Opus-xhigh model is the structurally-separated reviewer, never the author; §11.4.209 sets the review's MODEL + EFFORT, not its independence / iterate-until-GO (§11.4.134) / captured-evidence (§11.4.5/.69/.107). (C) **Honest boundary (§11.4.6)** — a review claimed done that ran on a weaker model / lower effort than Fable-xhigh (or Opus-xhigh on genuine Fable-unavailability) is a §11.4 bluff at the review-substrate layer; the review's captured evidence records which model + effort actually performed it (+ the Fable-unavailability FACT on fallback). STRENGTHENS §11.4.125/.142/.194/.134 — does NOT replace them; it pins the review's execution substrate. Classification: universal (§11.4.17). Composes §11.4.125/.134/.142/.194/.6/.20/.70/.196. Propagation gate `CM-COVENANT-114-209-PROPAGATION` (literal `11.4.209`) + recommended gate `CM-CODE-REVIEW-FABLE-XHIGH` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.209. Non-compliance is a release blocker. No escape hatch — no `--review-any-model`, `--review-lower-effort-OK`, `--skip-fable`, `--non-xhigh-review`, `--fallback-without-fable-unavailability-proof` flag.

§11.4.210 — Zero-loss request/prompt intake: every operator request/prompt MUST be mechanically CAPTURED, TRACKED, and PROCESSED — never skipped, ignored, avoided, or lost (User mandate, 2026-07-15). Verbatim intent: "every request/prompt we issue MUST ALWAYS be respected, taken into account, executed and processed! No avoiding, ignoring, skipping or any form of loss!" Every operator request/prompt MUST be (a) CAPTURED in the §11.4.208 request-history ledger MECHANICALLY at intake via a WIRED verified-runtime-state `UserPromptSubmit` auto-capture hook (NOT a manual step, NOT a "tracked follow-up" — §11.4.210 PROMOTES §11.4.208(D)'s hook from partial-follow-up to MANDATORY per the §11.4.205 enforced-not-advisory pattern; forensic FACT 2026-07-15: an entire release-drive session's prompts went un-recorded because the hook was left un-wired); (b) TRACKED → a §11.4.202 workable item / §11.4.197 completion, or an explicit evidence-backed no-op-with-reason; (c) PROCESSED — respected/executed, never dropped. Loss at ANY of the three stages is a §11.4 PASS-bluff at the requirements-intake layer of §11.4.197's severity class. Docs-chain-bound (§11.4.106), gated (§11.4.75 — hook verified-installed in the live hook path + content-matched, ledger fresh); the autonomous loop (§11.4.126/.87/.94/.97/.103) processes every captured request (done-condition never satisfied while any captured request is un-processed, §11.4.147(e)); honest §11.4.6 (UNKNOWN never invented). Classification: universal (§11.4.17). Composes §11.4.197/.202/.208/.205/.106/.126/.75/.6/.147(e). Propagation gate `CM-COVENANT-114-210-PROPAGATION` (literal `11.4.210`) + recommended gate `CM-REQUEST-CAPTURE-HOOK-WIRED` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.210. Non-compliance is a release blocker. No escape hatch — no `--skip-request-capture`, `--manual-ledger-OK`, `--hook-optional`, `--drop-unprocessed-request`, `--ledger-may-be-stale` flag.

§11.4.211 — Merge-conflict resolution during main→feature/product/flavor merges MUST run on the Fable model at xhigh effort (Opus xhigh fallback) (User mandate, 2026-07-16). Verbatim operator mandate: "During the merge of main branch to all feature/product/flavor branches of tracks 2 - 4, if conflicts happen they MUST BE resolved using Fable model with xhigh effort! If Fable is not available, then with Opus with xhigh!" Any merge-CONFLICT resolution performed while integrating the canonical trunk (`main` / `master`) into a feature / product / flavor branch — the §11.4.188 regular main→feature merge cadence, the §11.4.195 branch-taxonomy integration, the §11.4.41 merge-first pipeline, the §11.4.167(D) trunk-into-stream sync AND the §11.4.167(I) approval-gated back-merge — MUST be carried out on the **Fable model at `xhigh` effort**; the merge-resolution analogue of §11.4.209 (which pins the same model + effort for code-REVIEW), pinning it for the DECISION-MAKING of conflict resolution (which side to keep, how to UNION both sides without loss, marker removal, semantic reconciliation of divergent edits). **(A) Fallback (closed, ordered):** ONLY when Fable is genuinely unavailable (not configured / rate-limited / unreachable, established as FACT §11.4.6 — never a guessed unavailability) the conflict resolution runs on **Opus at `xhigh` effort**; a lower-effort or a non-Fable/non-Opus model is NOT an acceptable substrate for resolving merge conflicts, `xhigh` REQUIRED in both cases (a careless resolution silently drops a commit or commits a marker). **(B) Scope:** the operator named tracks 2–4 (`feature` / `product` / `flavor`); GENERALISES per §11.4.17 to EVERY controlled main→(feature|product|flavor) merge on EVERY track that hits a conflict. **(C) Safety preserved:** sets the MODEL + EFFORT of the conflict-resolution work only, BOUNDED BY

§11.4.113 (NEVER force-push) / §9 / §9.2 (no commit lost, hardlinked pre-op backup before any large/risky merge) / §11.4.41 step 3 (ZERO conflict markers committed, ZERO file silently dropped, union of both sides preserved) / §11.4.188(4) anti-bluff (post-merge smoke GREEN + marker-scan empty + no-lost-commit) / §11.4.84 (quiescent-only) / §11.4.37/§11.4.71 (fetch-first). **(D) Honest boundary (§11.4.6):** a merge whose conflicts were resolved on a weaker model / lower effort than Fable-xhigh (or Opus-xhigh on genuine Fable-unavailability) is a §11.4 bluff at the merge-resolution-substrate layer — the merge's captured evidence records which model + effort resolved the conflicts (and, on fallback, the FACT of Fable's unavailability). STRICT EXTENSION of §11.4.209 to merge-CONFLICT resolution; STRENGTHENS §11.4.188/§11.4.195/§11.4.41 by pinning the substrate of their conflict-resolution step. Classification: universal (§11.4.17). Composes §11.4.209/.188/.195/.41/.113/§9/§9.2/.6/.167(D)(I)/.84/.37/.71/.196. Propagation gate `CM-COVENANT-114-211-PROPAGATION` (literal `11.4.211`) + recommended gate `CM-MERGE-CONFLICT-FABLE-XHIGH` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.211. Non-compliance is a release blocker. No escape hatch — no `--merge-conflict-any-model`, `--skip-fable-merge-resolution`, `--non-xhigh-merge-conflict`, `--resolve-conflict-lower-effort-OK`, `--fallback-without-fable-unavailability-proof` flag.

§11.4.212 — Main README is the canonical starting-point / entry point for ALL project documentation — no doc may be an orphan unreachable from README (User mandate, 2026-07-16). Verbatim operator mandate: "Manual tests catalog must have linking all through the main README as the starting point of this and any other Project documentation! This fact and mandatory rule MUST be added into the constitution." The main README (`README.md` at the repository root) is the canonical STARTING POINT / entry point for ALL project documentation, for THIS project and EVERY project incorporating this constitution. Every project document in the §11.4.65 export scope — INCLUDING the manual-tests catalog (the §11.4.185 manual-QA testing catalog), the §11.4.153 per-feature Status / Status_Summary docs, the Issues / Fixed trackers + their summaries (§11.4.12 / §11.4.53 / §11.4.57), and every guide / research note / changelog / script companion doc — MUST be reachable / linked (DIRECTLY or TRANSITIVELY) from the main README. No project document may be an ORPHAN — a doc that no link-path from README reaches is a §11.4.212 violation. STRICT GENERALISATION of §11.4.57 (which mandated a `Tracked-`

`Items + Status Documents` doc-link section covering the tracker / status set) to EVERY §11.4.65-scope document, so the README is the single navigable root of the project's documentation tree. **(A) Reachability = a link path.** README → guide → sub-doc is fine (transitive); the requirement is that SOME path from README reaches every in-scope doc, not that every doc is linked directly at top level. **(B) Always-sync (§11.4.59 / §11.4.106 docs-chain).** A newly-added §11.4.65-scope doc that no README-reachable path links is a §11.4.212 violation the moment it lands; the manual-tests catalog specifically MUST be linked from the README (the operator's load-bearing example). **(C) Honest boundary (§11.4.6).** Reachability is a link-graph property (every doc FINDABLE from README), NOT a claim about each doc's CORRECTNESS or freshness (those stay §11.4.44 / §11.4.106 / the doc's own gates). **(D) Composition with §11.4.57.** §11.4.57's `Tracked-Items` section remains the mandated README landing block for the tracker / status set; §11.4.212 adds that EVERY OTHER in-scope doc is likewise reachable — §11.4.57 is the specialisation, §11.4.212 the generalisation. Classification: universal (§11.4.17) — the consuming project supplies its README path, its manual-tests-catalog path, and its §11.4.65 doc-set enumeration per §11.4.35. Composes §11.4.57 / §11.4.59 / §11.4.65 / §11.4.106 / §11.4.153 / §11.4.185 / §11.4.44 / §11.4.12 / §11.4.53. Propagation gate `CM-COVENANT-114-212-PROPAGATION` (literal `11.4.212`) + recommended gate `CM-README-DOC-ENTRYPOINT-COMPLETE` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.212. Non-compliance is a release blocker. No escape hatch — no `--orphan-doc-OK`, `--readme-not-entrypoint`, `--skip-readme-link`, `--doc-unreachable-OK` flag.

§11.4.213 — FEATURE research-scheduling directive: recognized via all §11.4.140 forms, SCHEDULES (never synchronously executes) a deep, enterprise-grade research + implementation-planning effort as a tracked workable item (operator-directed mandate, 2026-07-16). Compact summary: extends the §11.4.202 reporting-directive family with a fourth registered action `FEATURE` — recognized via ALL SIX §11.4.140 grammar forms (`FEATURE :: x` / `DEFAULT::FEATURE :: x` / `/FEATURE x` / `/DEFAULT::FEATURE x` / `FEATURE ---> x` / the single-colon `FEATURE: x`, registered-action-only so lowercase prose like "feature: ..." never expands and is never asked about, §11.4.6) — that SCHEDULES rather than synchronously executes a deep research + implementation-planning effort: the directive creates a Type=Task / Status=Queued workable item (§11.4.15/.16/.54) by INVOKING the SAME §11.4.202

`report_item.sh` engine (never a duplicate implementation of the item-creation/ DB-sync/tracker-push machinery — a second divergent tracker-push path would itself be a bluff-surface regression) and appends it to a durable `docs/requests/feature_queue.md` queue (mirroring `BACKGROUND`'s `docs/requests/background_queue.md` pattern) so a scheduled request is never dropped; the item's comprehensive description (§11.4.148/.171) embeds the FULL research mandate as its acceptance-defining work program — deep web research on best incorporation (§11.4.8/.99/.150), systematic-debug of every enumerated weak-spot/ gap/danger-zone with a risk-free rock-solid design per gap (§11.4.102), always enterprise-grade/bleeding-edge/innovative, exhaustive documentation down to lines-of-code + micro-POCs + diagrams + SQL + templates (§11.4.65/.73), reuse-first investigation of vasic-digital/HelixDevelopment components kept decoupled (§11.4.28/.74/.177), full test-type + Challenges + HelixQA planning (§11.4.27/.169), fine-grained phased/task/subtask breakdown, enterprise-scalability + max-performance planning, OpenDesign UI/UX wireframes where applicable (§11.4.162/.190), full CodeGraph integration (§11.4.78/.79/.80), and fully-synced workable items across the SQLite SSoT + every connected external tracker (§11.4.93/.95/.148 D5), honestly SKIPPING any absent tracker (§11.4.10) — NEVER a faked push. The multi-day research itself is EXECUTED by the standing autonomous loop (§11.4.87/.94/.97/.103/.126) when it claims the item, driven to a genuinely COMPLETED-and-wired or explicitly evidence-backed CLOSED terminal state under §11.4.197 — never left un-wired in the backlog; "scheduled" is never reported as "done" (§11.4.6). Anti-bluff four-layer coverage: pre-build gate `CM-FEATURE-DIRECTIVE` (FUNCTIONAL — sources + RUNS the §11.4.140 parser, never grep-only, §11.4.108) + paired §1.1 mutation + a real-DB end-to-end evidence trail (temp SQLite DB, no mocks, §11.4.27). Classification: universal (§11.4.17). Composes §11.4.140/.202/.197/.8/.99/.150/.102/.65/.73/.28/.74/.177/.27/.169/.162/.190/.78/.79/.80/.93/.95/.14 Propagation gate `CM-COVENANT-114-213-PROPAGATION` (literal `11.4.213`) + recommended gate `CM-FEATURE-DIRECTIVE` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.213. Non-compliance is a release blocker. No escape hatch — no `--feature-runs-synchronously`, `--skip-feature-queue`, `--feature-without-research-mandate`, `--reimplement-tracker-push`, `--hardcode-project-in-engine`, `--skip-item-creation` flag.

EXTENSIONS landed 2026-07-17 (existing anchors strengthened — full text in `Constitution.md`; research-derived from a consuming project's status-without-custody forensics). Forensic FACTs (genericised, all measured 2026-07-17): 576 tracked items, 182 claiming a done/ready status, **173 (95%) with NO registered regression guard** — an operator manually sampled 6 done-claiming items and found **6/6 broken**, the expected draw of that arithmetic, not bad luck; recorded reopen-per-recorded-fix **52%** (published elite <10%), itself an UNDERCOUNT because recurrences re-entered the tracker as NEW ids and statuses moved with ZERO history rows (one item: 4 Reopened events, 0 terminal events); the standing guard suite's exit code read ONLY its FAIL count — measured live on a release-candidate target: **61 registered guards → PASS 0 / FAIL 5 / PENDING 56 → exit 0 "not blocked"** — and the release-tag tooling consulted the guard suite **ZERO times**, so to the release gate "never verified" and "verified good" were the same colour; the project's own history selected the discriminator — every fix whose done-claim was preceded by an on-device RED→GREEN polarity flip that HELD recorded zero reopens, and every fix confirmed at source-green bounced. **No new anchor number was minted**: §11.4.115 and §11.4.146 already said what the operator mandates and did not bind because they are prose consulted by agents, not conditions checked by seams — the deliverable is the MECHANICAL BINDING (the §11.4.205 test applied: "if I documented this and nobody implemented the hook, would anything catch it?"). **§11.4.115(F) — machine-written verdict pairs + guard-viability:** RED and GREEN are harness-written verdict files, never prose — carrying the item id, guard identity, polarity, exit code, the target's artifact fingerprint READ FROM THE TARGET AT RUN TIME, iterations (≥3, §11.4.50), and evidence files whose CLASS matches the defect's layer (a source-grep transcript can never satisfy a user-visible/pixel/audio class — the wrong-layer oracle dies by construction); GREEN requires a prior RED with exit ≠ 0 on the PRE-fix artifact and a DIFFERENT fingerprint (identical fingerprints prove the fix was never deployed); a guard never observed FAILing on the genuinely-broken artifact is unvalidated instrumentation and mints no verdicts (a guard structurally unable to FAIL — or to PASS — satisfies nothing); the canonical §1.1 mutation for a landed fix is the fix-commit's REVERT, and a mutation whose diff only deletes the literal strings the gate greps for is a refused tautology. Recommended gate `CM-GUARD-VERDICT-MACHINE-WRITTEN`. **§11.4.135 (verdict-coverage extension) — ABSENCE of a verdict blocks exactly as a FAIL does:** the suite's exit semantics MUST incorporate COVERAGE, never only the FAIL count; the release seam refuses the tag while `uncovered = registered ∧`

`topology-present ∧ no-verdict-for-the-candidate-fingerprint` is non-empty OR any FAIL exists; the release-tag tooling MUST consult the suite/verdict store mechanically; REQUIRES-REBUILD PENDING is honest on intermediate artifacts and self-clearing on the candidate (re-run on the candidate, or the item leaves the release's done-set — both forced, neither a permanent exemption); topology-absent guards are exempt ONLY via a checked-in target-pool topology map and MUST be enumerated in the release changelog as honest gaps (§11.4.3/§11.4.118) — never silent; a naive "PENDING always blocks" is REJECTED as a §11.4.201 FAIL-bluff; pre-existing unguarded backlogs adopt via a ONE-TIME monotone-decrease ratchet snapshot so day one is green and the disease cannot spread. Recommended gate `CM-RELEASE-VERDICT-COVERAGE` (self-tested with golden-good / golden-bad / negative-control per §11.4.205/§11.4.107(10)). **§11.4.146(D3) — status-custody mechanical binding:** a done/ready status write is REFUSED without the file chain `registry row keyed by the item's EXACT stable id (checked-in alias map for legacy keys — a guard keyed to a non-item id fails the join and satisfies nothing) → executable registered guard → §11.4.115(F) RED+GREEN verdict pair → class-matched evidence` — prose appears nowhere in the chain; Seam A = the workable-items engine (§11.4.93) refuses the write (project-agnostic per §11.4.28, custody paths are consumer DATA per §11.4.35); Seam B = a FULL-TABLE sweep gate on every pre-build run (not a diff check — catches raw-DB writes that bypass Seam A); Seam C = the §11.4.135 release seam; §11.4.205 construction order — the executable sweep + its golden-bad meta-fixture (a temp SSoT overlay with a chain-less terminal item; the meta-test RUNS the sweep and requires FAIL, and asserts the sweep is WIRED) land in the SAME commit as any governance text citing them. Recommended gate `CM-STATUS-CUSTODY`. Honest boundary (§11.4.6), carried explicitly: the binding proves no status outruns its evidence and the REPORTED defect cannot silently return under the guard's conditions — it does NOT prove the feature bug-free (family dedup §11.4.186 remains), does NOT make a miscalibrated oracle honest (§11.4.107(10) goldens are necessary, not sufficient), does NOT replace the §11.4.185 manual-QA sufficiency gate, and makes target availability a HARD dependency of status progress BY DESIGN (a visible block replaces a silent pass). All Classification: universal (§11.4.17) — no hardware/vendor/project literal; consumers supply their registry, verdict-store, sweep, ratchet, and topology-map paths as DATA per §11.4.35. Propagation gates for the literals `11.4.115` / `11.4.135` / `11.4.146` are unchanged and stay GREEN; the three recommended mechanism gates + paired §1.1 mutations are gate-code = separate work item.