

GEMINI

**Helix Constitution – Universal
GEMINI.md**

Field	Value
Revision	36
Created	2026-06-09
Last modified	2026-07-26T11:48:35Z
Status	active
Status summary	Amendment of 2026-07-26 — the EFFORT-tiering + 5-field-label batch (operator IMPORTANT mandate 2026-07-26, verbatim: "Besides dynamic determining proper models to use, we MUST decide the effort as well and switch between the effort settings same as we do now for chosen model! Labeling MUST BE extended for this details as well. Example: (T1/main - claude4 - opus - high)."). §11.4.182 EXTENSION amended (in-place, no re-mint): the work-stream label form extends from (T<N>/<branch> - <alias> - <model>) to (T<N>/<branch> - <alias> - <model> - <effort>) , effort ∈ {low
Issues	none
Issues summary	—
Fixed	§11.4.140 + §11.4.141 mirrors
Fixed summary	Carrier created in lockstep with the Constitution.md §11.4.140 + §11.4.141 additions.
Continuation	—

How inheritance works

This GEMINI.md is the **Gemini CLI carrier** of the Helix Constitution submodule. It is one of the four canonical agent context carriers (`constitution/CLAUDE.md` , `constitution/AGENTS.md` , `constitution/QWEN.md` , `constitution/GEMINI.md`) — these ARE the canonical root per §11.4.35, NOT consumer extensions. The full universal rule set lives in `constitution/Constitution.md` ; the sibling carriers (`CLAUDE.md` / `AGENTS.md` / `QWEN.md`) carry the same content for their respective agents. A consuming project's own repo-root `GEMINI.md` (if any) opens with the inheritance pointer and adds only project-specific rules; THIS file is the canonical carrier and does not.

For agents driven from a consuming project, the project's repo-root carrier SHOULD start with a clearly-marked inheritance pointer:

```
## INHERITED FROM constitution/GEMINI.md
```

```
All rules in `constitution/GEMINI.md` (and the `constitution/Constitution.md`  
it references) apply unconditionally. The project-specific rules below  
extend them.
```

```
1 1 4 1 4 0
```

```
1
```

Action-prefix recognition (§ . . – LAYER–)

The block below is the **canonical LAYER-1 action-prefix recognition instruction**, embedded VERBATIM (identical text) into every agent context carrier the constitution maintains. That textual identity across all four carriers IS the LAYER-1 universality: Gemini CLI reads this block via this carrier and applies the prefix expansion itself, so the action-prefix system works on Gemini CLI with zero extra setup.

§11.4.140 — Universal action-prefix system (ACTION_NAME ::) (User mandate, 2026-06-09; GRAMMAR_ADDENDUM 2026-06-09; arrow form 2026-07-02). When a user prompt's FIRST non-blank line starts with a recognised action prefix, you MUST: (1) look the action token up in the action registry `constitution/actions/registry.yaml` (or `$HELIX_ACTION_REGISTRY`); (2) if it is a registered action, REPLACE the prefix with that action's `expansion` text and apply its `rules` ; (3) execute the remainder of the prompt under the expanded instruction. **Five EQUIVALENT forms** — same action, same expansion, same execution: (1) `ACTION_NAME :: <rest>` (bare `::`), (2) `PREFIX::ACTION_NAME :: <rest>` (namespaced `::`), (3) `/ACTION_NAME <rest>` (bare slash), (4) `/PREFIX::ACTION_NAME <rest>` (namespaced slash), (5) `ACTION_NAME ---> <rest>` (bare arrow) $\equiv$ `PREFIX::ACTION_NAME ---> <rest>` (namespaced arrow). Thus `BACKGROUND :: x` $\equiv$ `DEFAULT::BACKGROUND :: x` $\equiv$ `/BACKGROUND x` $\equiv$ `/DEFAULT::BACKGROUND x` $\equiv$ `BACKGROUND ---> x` $\equiv$ `DEFAULT::BACKGROUND ---> x` . `PREFIX` is an action NAMESPACE; the reserved default namespace is **DEFAULT** , and an action runs WITH or WITHOUT the prefix. Grammar (all hold): anchored at the FIRST non-blank line only (mid-prose tokens never match); the action token AND the namespace are UPPERCASE-only `[A-Z][A-Z0-9_]*` (lowercase never matches); the namespace separator `::` inside the token carries NO surrounding spaces (`PREFIX::ACTION_NAME`), DISTINCT from the action-body separator `" :: "` (one ASCII space on each side of `::` — avoids `C+ + Foo::Bar` , YAML `key: value` , URLs) in forms 1/2, the arrow-body separator `" ---> "` (one ASCII space on each side) in form 5, and the slash-body separator (one space) in forms 3/4; stacked prefixes (`A :: B :: rest`) apply outer-to-inner, left-to-right (expand `A` , re-scan, expand `B` , then the residual is the task); a leading `\` escapes the prefix for the `::` , arrow, and slash forms (`\BACKGROUND :: x` , `\BACKGROUND ---> x` , `\ /BACKGROUND x` — treat literally, strip the backslash, NO expansion) so action names can be discussed. **Conflict rule (slash form):** `/ACTION_NAME` (form 3) is honored as the action ONLY when `ACTION_NAME` (case-folded) does not collide with a built-in/host slash command (registry `slash_bare: auto` + `slash_conflicts: [..]`); form 4 (`/PREFIX::ACTION_NAME`) is ALWAYS unambiguous and always honored. An unknown token that matches the grammar shape (any of the 5 forms) but is NOT registered is NEVER silently

expanded or silently dropped — ask which registered action was meant (§11.4.66 / §11.4.105) or treat it as a literal prompt, NEVER invent an expansion (§11.4.6); any prompt not satisfying the grammar is an ordinary prompt and the system is a no-op. The registered action **BACKGROUND** expands to: *"The following prompt that we will provide MUST BE executed in background in parallel with all main work streams using the subagents-driven development approach! All work done MUST PRODUCE rock solid evidence covered with hard physical proof(s) that all done is working as expected and as specified without any false results and without any bluff!"* (composes §11.4.20 / §11.4.70 subagent-driven, §11.4.58 / §11.4.103 parallel streams, §11.4.89 background execution, §11.4.5 / §11.4.69 / §11.4.107 captured physical evidence, §11.4 anti-bluff). A second registered action **REMINDER** re-surfaces previously-scheduled, CRITICAL, status-UNCERTAIN work: FIRST verify the ACTUAL current status from captured evidence (never assume done or not-done — §11.4.6 no-guessing), THEN act on the delta — report the captured proof if genuinely complete, resume from the exact point if partial, action it NOW if not started, or surface the block per §11.4.66/ §11.4.101 — always producing a status verdict (composes §11.4.6 / §11.4.87 / §11.4.94 / §11.4.97 / §11.4.103 / §11.4.108 / §11.4.130 / §11.4.147).

Severity / handling-priority markers (§11.4.140 extension, 2026-07-16).

Three further registered actions tag the REMAINDER of the prompt with a HANDLING PRIORITY (as **BACKGROUND** tags it "run in background"):

CRITICAL (highest-priority / potentially release-blocking — maximum urgency + rigor ahead of lower-priority work, track it, auto-activate §11.4.102 systematic-debugging when an issue is involved, do NOT defer it),

IMPORTANT (high-priority, above routine work but below CRITICAL — track it, full anti-bluff rigor), and **NOTE** (capture the remainder as durable

CONTEXT — request-history §11.4.208 / §11.4.210 + persistent memory when a non-obvious project fact — applied when relevant, NOT an urgent action unless it contains an explicit action; §11.4.6 record only what the note says, never invent). All six grammar forms work (**CRITICAL**: x ≡

CRITICAL :: x ≡ **CRITICAL** ---> x ≡ /**CRITICAL** x ≡ /

DEFAULT::**CRITICAL** x). Registering **NOTE** promotes **NOTE**: from an ordinary-prose NO-OP to a registered action (ordinary-prose single-colon examples are now **TODO**: / **WARNING**: / **FIXME**:).

The system is UNIVERSAL (every CLI agent reads this block via its context carrier per §11.4.35), extensible (new action = new registry row), decoupled + reusable (§11.4.28), and loads out-of-the-box. Classification: universal (§11.4.17). **Canonical authority:** constitution submodule `Constitution.md` §11.4.140. Non-compliance is a release blocker. No escape hatch — no `--skip-action-prefix`, `--ignore-prefix`, `--no-registry`, `--invent-expansion-OK`, `--single-layer-only` flag.

1 1 4 1 4 1

Token-efficiency mandate (§ . .)

§11.4.141 — Token-efficiency mandate (research-derived + operator mandate, 2026-06-09). Every project worked on by AI coding agents **MUST** cut token spend (input AND output) toward **30–40% of current (a 60–70% reduction)** WITHOUT degrading quality/performance/safety or breaking any existing mechanism, via a composable, safety-ranked measure set: (1) **prompt-cache the static governance prefix** — the always-loaded governance forms a byte-stable cache-breakpointed prefix with no volatile bytes ahead of it; cache reads cost $\sim 0.1\times$ base input (the dominant cost driver — measured $\sim 170\text{K}$ tokens of governance re-sent every turn, externally corroborated by Claude Code issue #24147); caching is transparent so it removes no rule, weakens no gate, changes no verdict — only billing (PRIMARY, biggest + safest lever); the operative discipline is to **batch governance edits and keep CLAUDE.md + the constitution byte-stable mid-session** — every governance edit busts the prompt cache (BASELINE.md finding); (2) **subagent model-tiering + output-to-file** — mechanical non-judgment work (search/grep/status/doc-export/read-only probes) to a Haiku-class model, the strong model RESERVED for all reasoning/verdicts/fix-design (§11.4.102)/code-review (§11.4.125)/demotion (§11.4.7), large output persisted to a file not an inline 350–520K-token transcript; the cheap model never emits a PASS so §11.4.50 + anti-bluff are untouched; (3) **thin always-loaded INDEX + on-demand detail** — concise index (one line per fix/anchor, EACH carrying the literal `11.4.N` token so propagation gates pass) with the canonical full text kept gate-scanned in `constitution/Constitution.md` and reachable in one hop — a de-duplication realising §11.4.35, never a deletion; (4) **CodeGraph/retrieval-first over full-file loading** (§11.4.78/§11.4.79/§11.4.80); (5) **output-token reduction** — terse status

+ `effort:"low"` on the mechanical allowlist only; (6) **tool-call batching + no re-reads**; (7) **compaction/context-editing for long sessions**. **Mandatory measured proof:** a token-accounting harness measures tokens-per-development-cycle BEFORE vs AFTER on a frozen deterministic workload from the authoritative `usage` object (input/cache_read/cache_creation/output split; NEVER `tiktoken`, NEVER the client-side cost estimate), reproduced N times (§11.4.50), `pass = AFTER ≤ 40% of BEFORE` OR the measured best-safe reduction with a cited cold-cache reason; the AFTER run MUST show ZERO regression on the pre-build sweep + meta-test mutation sweep + propagation gates + a strong-model reasoning probe + a cache-warm proof (`cache_read_input_tokens > 0`) — cost reduction with quality regression is a §11.4 FAIL. The headline number is the *measured* reduction, never the design estimate (§11.4.6/§11.4.123). No measure may break/degrade any existing mechanism, and the rule is structured so none can. Composes

§11.4.5/.6/.20/.40/.50/.58/.69/.70/.78/.79/.80/.103/.106/.123/.125/.128/§12.6/§1.1. Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-141-PROPAGATION` (literal `11.4.141`) + recommended gate `CM-TOKEN-EFFICIENCY` + paired §1.1 mutation (inject a pre-breakpoint volatile token → cache collapses → measured reduction falls below bar → gate FAILs). **Canonical authority:** constitution submodule `Constitution.md` §11.4.141. Non-compliance is a release blocker. No escape hatch — no `--skip-token-efficiency`, `--no-cache-governance`, `--assert-reduction-without-measuring`, `--tier-down-reasoning`, `--inline-all-governance`, `--tiktoken-estimate-OK` flag.

§11.4.142 — Universal code-review mandate — every change reviewed, always, no exception (User mandate, 2026-06-09). Verbatim operator mandate: "ALL changes we do MUST pass through the code review step!!! ALWAYS!!!" EVERY change made to ANY repository this Constitution governs — without exception — MUST pass through an INDEPENDENT code-review step BEFORE it is accepted, committed, or built. NO change class is exempt: source code, fixes, tests, gates, meta-test mutations, documentation, doc-tooling, build/CI scripts, configuration, governance files (Constitution / CLAUDE / AGENTS / QWEN), conductor main-stream edits, sub-agent output, refactors, single-line edits — if a diff exists, it gets an independent review. This is the ABSOLUTE form of §11.4.125 (code-review agent gate after a batch, before pre-build sweep + main build): §11.4.142 strips every implicit scoping §11.4.125's "after all fixes/changes/ implementations are done" phrasing could leave open — no "just a doc edit", no

"just a one-liner", no "the author already self-reviewed (§11.4.92)", no "trivial change" carve-out. **Independence is load-bearing** — the reviewer MUST be structurally separated from the author (a dedicated code-review agent, subagent-driven by default per §11.4.70 / §11.4.20, or a distinct human), NEVER the author re-reading their own work; §11.4.92's multi-pass self-evaluation is the AUTHOR-side discipline and PRECEDES (never satisfies) §11.4.142. **The review is itself anti-bluff** (§11.4 / §11.4.1) — a rubber-stamp "looks good" is not a review; it MUST genuinely analyse correctness, safety (no host §12 / data §9 / target-hardware §11.4.133 regression), will-it-really-work (no solve-A-create-B), end-user behaviour (§11.4 / §107), and test-genuineness (§1.1), its conclusions captured evidence per §11.4.5 / §11.4.69, and **it iterates to a clean GO per §11.4.134** — any finding (BLOCKING / nit / warning) re-arms the loop, acceptance only on ZERO new findings + ZERO warnings with rock-solid physical evidence. Honest boundary (§11.4.6): a passing review does NOT replace §11.4.108 four-layer runtime-signature verification nor the §11.4.40 full-suite retest — it is one of MULTIPLE STRONG LAYERS, and the FIRST one every change crosses. Composes §11.4.1 / §11.4.4 / §11.4.5 / §11.4.6 / §11.4.20 / §11.4.40 / §11.4.69 / §11.4.70 / §11.4.92 / §11.4.108 / §11.4.110 / §11.4.125 / §11.4.134 / §107 / §1.1. Classification: universal (§11.4.17) — the consuming project supplies its reviewer-dispatch mechanism + change-acceptance seam (commit wrapper / merge queue / PR gate) per §11.4.35. Propagation gate `CM-COVENANT-114-142-PROPAGATION` (literal `11.4.142`) + recommended gate `CM-EVERY-CHANGE-REVIEWED` (every accepted change carries a fresh independent-review-completed marker for its diff, produced by a reviewer distinct from the author, before acceptance/commit/build) + paired §1.1 mutation (strip the literal → propagation gate FAILs; accept a diff with no independent-review marker → `CM-EVERY-CHANGE-REVIEWED` FAILs). **Canonical authority:** constitution submodule `Constitution.md` §11.4.142. Non-compliance is a release blocker. No escape hatch — no `--skip-review`, `--no-review`, `--trivial-change-exempt`, `--doc-edit-exempt`, `--self-review-suffices`, `--review-after-commit` flag.

§11.4.143 — Real-user-journey mandate for video-streaming-app full-automation tests (User mandate, 2026-06-10). Verbatim operator mandate: "All video streaming apps ... require to choose some title and to press proper UI button to start or resume playing! Proper UI interaction to play exact show with proper content and subtitles is MANDATORY! Without it we just eventually play on 2nd display sample, and that's it mostly! THIS MUST BE ADDED as MANDATORY

RULE regarding testing of any video streaming app in general with full automation tests! Universal in root constitution + a consuming project's extensions." Any full-automation test that asserts a video player / streaming application plays content MUST drive the REAL end-user journey through the app's OWN UI — launch → BROWSE the actual catalog → choose a SPECIFIC title → press the real Play/Resume button → confirm THAT chosen content is genuinely playing, with its correct subtitles, on the intended routing target. A test that bypasses the journey with a sample / demo / built-in-loop clip, a deep-link / `am start -a VIEW` / intent shortcut, a synthetic or pre-staged stream, or any path that does NOT exercise the app's own browse-select-play UI is a §11.4 PASS-bluff at the user-journey layer: it validates ROUTING (that *something* reaches the display) while leaving the user-visible behaviour the operator cares about — "the show I picked actually plays" — unproven (the operator's "we just eventually play on 2nd display sample, and that's it mostly" gap). The mandate (ALL hold): (1) **Real journey, not a shortcut** — launch → catalog browse → specific-title selection → real Play/Resume press through the app's own UI (§11.4.48 UI-driven), NEVER a deep-link / intent / sample / loop-clip shortcut; (2) **Chosen-content confirmation** — the PASS proves the SPECIFIC selected title is playing (not merely that pixels move on the target) via the §11.4.107 liveness battery on the §11.4.136 real-content path + the §11.4.137 subtitle content-correctness oracle for the chosen title's captions; (3) **Non-introspectable UIs use the pixel oracle** — when the app's accessibility hierarchy is blank/unreliable (TV-Compose / leanback / canvas / GL), DRIVE input + ASSERT content via the §11.4.117 CV/OCR pixel oracle, never a hierarchy-only tool; (4) **Login via the credential single-source** (§11.4.10), never hardcoded, never logged; (5) **Honest SKIP, never a faked PASS** — where the autonomous journey is genuinely infeasible (hard human-only login / CAPTCHA, geo-block per §11.4.3, secure-surface pixel-blanking per §11.4.112) the test is `operator_attended SKIP-with-reason` per §11.4.52 + §11.4.3 with a tracked migration item — NEVER a metadata-only / sample-played / routing-only PASS. Honest boundary (§11.4.6): "the routing fired so the title is playing" is a guess — only the chosen-content liveness + subtitle oracle on the real journey proves it. Composes §11.4.48 / §11.4.107 / §11.4.117 / §11.4.136 / §11.4.137 / §11.4.52 / §11.4.3 / §107 / §1.1. Classification: universal (§11.4.17) — the consuming project supplies its concrete app roster, login-credential source, routing target, and UI-driving / pixel-oracle harness per §11.4.35. Propagation gate `CM-COVENANT-114-143-PROPAGATION` (literal `11.4.143`) + recommended gate `CM-VIDEO-REAL-JOURNEY-TEST` (every video-streaming-app playback test drives the

real browse-select-play journey + confirms the chosen content + subtitles, or SKIPS-with-reason) + paired §1.1 mutation (replace a real-journey test's browse-select-play path with a deep-link / sample shortcut → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.143. Non-compliance is a release blocker. No escape hatch — no `--sample-playback-ok` , `--skip-real-journey` , `--deep-link-suffices` , `--routing-only-pass` , `--no-title-selection` flag.

§11.4.144 — Tracked/recorded-device availability-following mandate (User mandate, 2026-06-10). Direct operator mandate: every device the project is tracking / following / recording (test / debug / manual-testing device, across every reachable transport — USB / wireless ADB / SSH / serial / network introspection API) MUST be availability-FOLLOWED — its connection state continuously monitored, any drop handled, never silently abandoned and never presented as a continuous recording. The §11.4.128 always-on recorder KNOWS a tracked device is absent (its per-device loop guards on the reachable state) but, lacking a following discipline, merely spins idle — no data captured, no offline event logged, no resume, no escalation — so the recording corpus presents a continuous timeline with a silent hole, a §11.4 PASS-bluff at the recording-integrity layer (it claims continuous capture while no data exists for the gap). On a tracked device leaving its reachable state the system MUST, automatically: (a) **DETECT** the drop and **log an honest offline event** into the recording corpus (§11.4.6 — a silent gap presented as continuous capture is a fabricated-continuity bluff; §11.4.128 — a silent recording gap defeats always-on recording); (b) **WAIT** for the device to return using the project's ALREADY-DEFINED reconnection timings — never invented numbers (§11.4.6), the SAME grace / reconnect / poll budgets the project's recovery path already uses; (c) **RE-ATTACH** — resume recording / tracking the moment the device returns and log an honest online / resume event; (d) **ESCALATE** to the project's sanctioned device-recovery path (per §11.4.69 feature class `device_recovery`) if the device does not return within the defined timeout — through the sanctioned, authorization-gated recovery entry point ONLY, never bypassing its gate and never performing a destructive recovery (e.g. a power-cycle) autonomously without that authorization (§11.4.21 / §11.4.101 — high-blast-radius recovery is gated; while blocked the system keeps following the device, logging the blocked-escalation honestly). A tracked-device drop that produces a silent corpus hole, a never-resumed recording, or a never-escalated permanent absence is the

bluff this anchor forbids. Composes §11.4.128 (always-on recording — §11.4.144 closes its drop-handling gap) / §11.4.69 (`device_recovery` sink-side positive evidence) / §11.4.6 (honest offline / online events, reused-not-invented timings) / §11.4.14 (watchdog children reaped on stop) / §11.4.21 + §11.4.101 (gated, non-autonomous escalation). Classification: universal (§11.4.17) — the consuming project supplies its concrete tracking transport, device set, already-defined reconnection timings, and sanctioned recovery entry point per §11.4.35.

Propagation gate `CM-COVENANT-114-144-PROPAGATION` (literal `11.4.144`) + recommended gate `CM-DEVICE-AVAILABILITY-FOLLOWED` + paired §1.1 meta-test mutation (strip the watchdog wiring / let an absence go unlogged → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.144. Non-compliance is a release blocker. No escape hatch — no `--skip-availability-following` , `--silent-recording-gap-OK` , `--no-reconnect-wait` , `--invent-reconnect-timing` , `--skip-recovery-escalation` , `--autonomous-power-cycle-OK` flag.

§11.4.145 — Independent multi-angle impact-research per change (User mandate, 2026-06-10). For EVERY fix / change / new feature, INDEPENDENT impact-research agents (subagent-driven §11.4.70/§11.4.20, structurally separate from the author — §11.4.92 self-eval PRECEDES, never satisfies — adversarial "refute-the-change" stance to defeat the documented LLM confirmation-bias failure mode) MUST research the change AND its connected/dependent features across a CLOSED SET OF EIGHT ANGLES — (1) correctness/logic, (2) regression (what existing working feature could break — via call-graph impact enumerating every direct + transitive caller and contract-dependent feature, each mapped to its test), (3) latent/dangerous-code — the recents class (runtime-only failure, interface/ABI/contract mismatch, concurrency/data-race, resource lifecycle), (4) security (taint/injection/secret/unsafe-API), (5) performance (latency/memory/thermal under load vs baseline), (6) host/data/target-hardware safety per §12/§9/§11.4.133, (7) cross-feature interaction (shared state/timing/hardware contention), (8) business-logic conformance (matches the spec AND the connected features' contracts, not merely compiles). Forensic anchor (FACT, project-internal): the recents / 3-button-nav path shipped a latent AIDL-interface-contract mismatch that PASSED code review yet failed at RUNTIME ONLY (ANR on a recents-reaping gesture) because no review angle interrogated the interface contract / concurrency-lifecycle / silently-broken connected feature — the §11.4.108 SOURCE → ARTIFACT → RUNTIME → USER-

VISIBLE gap caught only after the fact; §11.4.145 shifts that discovery LEFT. Each angle MUST name the TOOL(S) used + obtain the required DEPTH (the diff, the full changed unit, its declared contract, the spec section, every connected feature — NEVER the diff lines alone) + emit a captured-evidence conclusion per §11.4.5/§11.4.69 ("looks fine" forbidden, §11.4.1); a genuinely-N/A angle is recorded NOT_APPLICABLE: <reason> per §11.4.6, never silently skipped. Output = a per-change impact-research REPORT (one section per angle: tool + evidence path + verdict) + a single GO/NO-GO that BLOCKS acceptance/commit/build on ANY unmitigated risk in ANY angle; the change is fixed/mitigated + affected angles re-researched, iterating to a CLEAN GO (zero unmitigated risk + zero warning) per §11.4.134 before the §11.4.125 review gate. Honest boundary (§11.4.6): a GO proves cross-angle internal-consistency + bounded blast-radius — it does NOT replace §11.4.108 runtime-signature verification nor §11.4.40 full-suite retest; it is one of MULTIPLE STRONG LAYERS, the FIRST research pass every change crosses. Composes/strengthens §11.4.92/.125/.108/.110/.134/.78/.79/.107/.85/.133/§12/§9/.99/.6/§1.1. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-145-PROPAGATION (literal 11.4.145) + recommended gate CM-IMPACT-RESEARCH-PER-CHANGE + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; accept a change with a report missing an angle or an unmitigated NO-GO → CM-IMPACT-RESEARCH-PER-CHANGE FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule Constitution.md §11.4.145. Non-compliance is a release blocker. No escape hatch — no --skip-impact-research, --single-angle-suffices, --author-researches-own-change, --diff-skim-OK, --no-go-overrideable, --trivial-change-no-research flag.

§11.4.146 — Reproduce-first test + same-test-confirms-fix + mandatory extend-to-all-cases workflow (User mandate, 2026-06-10). Every reported problem MUST be handled by a NAMED three-step test workflow — §11.4.146 does NOT re-author its component disciplines, it NAMES + BINDS them into the operator's exact per-defect sequence and adds ONLY the two emphases the components leave implicit. **(STEP 1 — REPRODUCE-FIRST + INVESTIGATE)** before any fix, author the §11.4.43 RED test as a §11.4.115 RED-baseline-on-the-broken-artifact (reproduce the defect on the CURRENT pre-fix artifact, capture defect-present physical evidence per §11.4.5/§11.4.69/§11.4.107); NEW EMPHASIS (D1) that same RED test is ALSO a deliberate investigation instrument per §11.4.102(A) — it MUST gather ADDITIONAL forensic data characterising the

defect (triggers, boundaries, input/topology scope, adjacent failure modes) that feeds BOTH the fix design AND the STEP-3 extend scope; a RED test used only as a binary present/absent gate with no characterisation satisfies §11.4.115 but NOT §11.4.146 STEP 1. **(STEP 2 — SAME-TEST-CONFIRMS-FIX)** after the fix the SAME test source (the §11.4.115 polarity switch flipped `RED_MODE=1→0`) confirms the defect ABSENT — RED-on-broken then GREEN-on-fixed, both captured, on a clean target per §11.4.108/§11.4.139, validated-first-after-redeploy per §11.4.130, deterministic per §11.4.50; no separate happy-path test substitutes (the §11.4.43/§11.4.115 PASS-bluff). **(STEP 3 — EXTEND-TO-ALL-CASES, mandatory per-fix)** NEW EMPHASIS (D2) immediately after STEP 2 the test MUST be FANNED OUT across the full case-space of the SAME functionality — all flows, valid + invalid + boundary (§11.4.85 empty/max/off-by-one) + concurrent (§11.4.85 contention) + failure-injection (§11.4.85 chaos) + topology variants (§11.4.3) — confirming no issue of any kind, with PROVABLE enumerated coverage per §11.4.118 (a listed case-set with per-case outcome, never "no other issues found"); a REQUIRED step, NOT deferred to a release-cycle discovery sweep and NOT reduced to a single guard; each user-visible case carries rock-solid physical evidence per §11.4.123 + is registered into the §11.4.135 standing regression-guard suite; a newly-discovered fan-out defect triggers §11.4.4 test-interrupt + re-enters STEP 1. (ANTI-BLUFF, all steps) every PASS is rock-solid CAPTURED physical evidence per §11.4.123 (§11.4.5/§11.4.69/§11.4.107) — metadata-only / config-only / absence-of-error / grep-without-runtime PASS forbidden (§11.4/§11.4.1); unclear validation method ⇒ deep-research-before-declaring-untestable per §11.4.123/§11.4.8/§11.4.99; an operator-found defect the green suite missed triggers §11.4.138. Honest boundary (§11.4.6): STEP 3 reduces the unknown-unknown surface but does not prove zero remaining defects (§11.4.118) — un-exercised cases stated as honest gaps, never silently implied clean. Composes §11.4.43/.115/.102/.130/.108/.139/.50/.85/.118/.135/.3/.123/.5/.69/.107/.138/.4/§107/§1.1. Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-146-PROPAGATION` (literal `11.4.146`) + recommended gate `CM-REPRODUCE-FIRST-THEN-EXTEND` (every closed defect's fix carries a §11.4.115 reproduce-first polarity test with captured defect-characterisation [STEP 1] + its `RED_MODE=0` GREEN confirmation [STEP 2] + an enumerated per-functionality extend case-set registered into the §11.4.135 suite [STEP 3]; a closure with the reproduce → confirm pair but NO enumerated extend case-set FAILs) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; close a defect with only the reproduce → confirm pair and no extend-case-set →

CM-REPRODUCE-FIRST-THEN-EXTEND FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md §11.4.146. Non-compliance is a release blocker. No escape hatch — no

--skip-reproduce-first , --fix-without-red , --skip-extend-to-all-cases , --reproduce-confirm-suffices ,
--defer-extend-to-release-sweep , --single-guard-suffices , --no-defect-characterisation flag.

§11.4.147 — Crashed-agent respawn-until-complete + no-work-loss registry mandate (User mandate, 2026-06-10).

Verbatim operator intent: "any agent that crashed because of something MUST BE respawned and finish its work at some point! We MUST NOT lose any work, forget about or have it corrupted!" Forensic case study (FACT): dispatched subagents died on TRANSIENT causes (API Error: Server is temporarily limiting requests rate-limit killed 5 at once, API Error: socket connection closed unexpectedly killed 2, one left PARTIAL edits — two source files written, the dependent SystemUI sliders + doc + test NOT done), each re-dispatched by pure conductor vigilance with NO mechanical guarantee — the moment conductor context is lost / compacted / the conductor itself crashes, an in-flight-but-dead work unit is silently dropped (lost / forgotten / corrupted). Every dispatched agent/subagent MUST be tracked through its full lifecycle so a crash NEVER loses, forgets, or corrupts its work; a crashed agent is NOT a completed agent (§11.4.6 — "it probably finished before dying" is a forbidden guess; abnormal termination is positive evidence the unit is OPEN). The mandate (ALL hold): **(a) durable agent REGISTRY** — every dispatched agent has an append-only machine-readable registry entry (stable id, task + §11.4.58 file-scope, declared output path(s) + §11.4.5/§11.4.69 evidence dir, dispatch timestamp, status from the CLOSED SET {dispatched | in-flight | crashed | respawned | complete}), reusing the §11.4.116 sync substrate (JSONL event stream + atomic status snapshot; "an entry with no dispatch-event cannot show complete "); the registry is the SINGLE SOURCE OF TRUTH for "is this work owed?" and an unregistered dispatch is itself a violation; **(b) mechanical CRASH-DETECTION + RESPAWN-until-complete** — any abnormal termination (transient rate-limit/socket-close, any non-completion exit, or a terminal error after the runtime's own retries are exhausted) flips the entry to crashed + keeps the unit OPEN, and the unit is respawned (fresh agent claims the same id + scope + preserved state) and re-respawned until an agent reaches complete ; respawn is the safe/reversible/bounded decision taken autonomously per §11.4.101 (never

blocks the loop for a human), transient causes use the project's ALREADY-DEFINED backoff budgets (never invented, §11.4.6), a deterministically-reproducible non-transient terminal error is investigated per §11.4.102 before further respawn (never a blind retry loop); **(c) PARTIAL-STATE preserve** → **§11.4.84-check** → **resume-or-clean-restart** — the crashed agent's uncommitted edits + output doc are PRESERVED (never silently discarded → lost, never blindly committed → corruption), the respawn runs the §11.4.84 quiescence check on the preserved tree (every modified file accounted-for vs scope, no mutation/ // always pass / `_mutated_*` residue, no half-written/torn artifact), then EITHER RESUMES idempotently when it passes OR CLEAN-RESTARTS from a known-good base (§9.2 pre-op backup, reversible §11.4.101) when inconsistent — nothing lost, nothing corrupted; per-agent `git worktree` isolation (§11.4.58 L4 / §11.4.84) keeps the partial tree from contaminating other streams; **(d) "crash ≠ done"** **COMPLETION criterion** — a unit is DONE only when an agent reaches `complete` with its required captured evidence/output landed (§11.4.5/§11.4.69 + the §11.4.116 verdict-carries-evidence-path rule); the endless-loop done-condition (§11.4.87/§11.4.94/§11.4.97/§11.4.126) MUST NOT read satisfied while any entry is `dispatched / in-flight / crashed / respawned -not-yet- complete`, the zero-idle survey (§11.4.94) treats every non-`complete` entry as an OPEN item to reclaim, and a registry showing `complete` without landed evidence is a §11.4 PASS-bluff at the agent-lifecycle layer. Honest boundary (§11.4.6): the registry + respawn guarantee work is not lost/corrupted, NOT that it is correct — the respawned output still crosses §11.4.108 + §11.4.125/§11.4.142 review + §11.4.40 retest; §11.4.147 is the durability layer beneath those, the agent-side analogue of §11.4.144 (same detect → wait/backoff → re-attach/respawn → escalate shape). Composes §11.4.6/.58/.84/.87/.94/.97/.101/.116/.102/.108/.125/.142/.40/.128/.144/ §9.2/§1.1. Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-147-PROPAGATION` (literal `11.4.147`) + recommended gate `CM-CRASHED-AGENT-RESPAWN-TRACKED` + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; mark a `crashed` entry as the loop's done-condition `complete` without landed evidence, OR drop a dispatched agent without a registry entry → `CM-CRASHED-AGENT-RESPAWN-TRACKED` FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.147. Non-compliance is a release blocker. No escape

hatch — no --skip-agent-registry , --crash-equals-done , --no-respawn ,
--discard-partial-state , --blind-commit-partial , --forget-dead-agent ,
--loop-done-with-crashed-entries flag.

§11.4.148 — Workable-item integrity (status+type+id) + comprehensive structured description + bidirectional external-tracker sync + BLOCKED unblock-choices mandate (User mandate, 2026-06-10). BINDS + STRENGTHENS the workable-item discipline (§11.4.15 status / §11.4.16 type / §11.4.54 id / §11.4.21 Operator-blocked / §11.4.91 description-clarity / §11.4.93 + §11.4.95 SQLite SSoT / §11.4.106 docs_chain) into ONE integrity contract spanning DB ↔ docs ↔ external tracker, adding the operator's three emphases: **(D1) no item without a valid status + valid type + stable id — on ALL three surfaces** (an item missing any of the three in the DB, the rendered docs, OR the external tracker FAILs the validator — release-blocker; the id is the cross-surface binding key); **(D2) comprehensive structured description per item** — WHAT it is (§11.4.91 ≥6-word/≥40-char clear meaning) + HOW it manifests + HOW to reproduce (§11.4.115/.146) + ACCEPTANCE CRITERIA (§11.4.123/.69 captured-evidence verdict), in the §11.4.93 DB description column AND the docs AND the tracker, never a stub/§11.4.91-anti-pattern fragment; **(D3) BLOCKED items carry WHY + enumerated unblock CHOICES** — tightens §11.4.21: every Operator-blocked item (incl. Blocked / BLOCKED documented alias normalised to canonical Operator-blocked , never a silent fork §11.4.6) MUST enumerate the closed list of decisions/actions that would unblock it ([A]...·[B]...·[C]... , mirroring §11.4.66's 2–4-option shape); a BLOCKED item with no enumerated choices FAILs; **(D4) regular never-missed bidirectional DB ↔ docs ↔ tracker sync** — the §11.4.93/.95 git-tracked SQLite DB is the SSoT, docs + tracker DERIVED; full sync (db-to-md / md-to-db byte-identical round-trip §11.4.93 + tracker push) runs regularly + on every change, §11.4.86 drift-proof fingerprint (sha256 of the sorted item keyset, NOT mtime) gates freshness, §11.4.106 docs_chain-bound so drift is mechanically caught not vigilance-dependent; **(D5) generic external-tracker push** carries statuses (collapsed onto the tracker's native set per §11.4.33/.112 when fixed, precise value preserved in a header, never lost), types, assignee (§11.4.104 participant handle, UNSET defaults from a project env var, never hardcoded/logged §11.4.10), and sub-tasks, IDEMPOTENT (dry-run-then-real, match-by-stable-key [<ID>] prefix / id custom-field, present⇒UPDATE/absent⇒CREATE, rate-limited, credential-redacted, sink-side created=N updated=M failed=0 proof §11.4.69), the MACHINERY project-agnostic §11.4.28

(consumer registers its tracker/list/field-map at runtime). Anti-bluff §11.4: every sync + validator pass carries captured evidence §11.4.5/.69; honest boundary §11.4.6 — guarantees well-formed items + agreeing surfaces, NOT that the underlying work is correct (still crosses §11.4.108/.40/.123). Composes

§11.4.15/.16/.21/.33/.34/.54/.66/.86/.91/.93/.95/.104/.106/.112/.123/.10/.28/.69/.6/

§1.1. Classification: universal (§11.4.17) — consumer supplies its DB path, id prefix, tracker service + list/board id + field map + default-assignee env var, docs_chain context per §11.4.35. Propagation gate CM-COVENANT-114-148-PROPAGATION (literal 11.4.148) + recommended gates CM-ITEM-INTEGRITY-STATUS-TYPE-ID / CM-ITEM-COMPREHENSIVE-DESCRIPTION / CM-BLOCKED-UNBLOCK-CHOICES / CM-TRACKER-SYNC-IDEMPOTENT + paired §1.1 mutation (gate-code = separate work item). **Canonical authority:** constitution submodule

Constitution.md §11.4.148. Non-compliance is a release blocker. No escape hatch — no --item-without-status , --item-without-type , --item-without-id , --stub-description-OK , --blocked-without-choices , --skip-tracker-sync , --one-way-sync-OK , --non-idempotent-tracker-push , --hardcode-assignee flag.

§11.4.149 — Per-workable-item testing-diary mandate (User mandate, 2026-06-10). Every workable item MUST carry an append-only TESTING DIARY — a chronological record of every test run against it — distinct from §11.4.93

item_history (lifecycle STATE transitions) + §11.4.55 Reopens (reopen cycles): the diary records TEST EXECUTIONS (which may or may not change status). The mandate (ALL hold): **(a)** a test_diary table additive to the §11.4.93/.95 SQLite SSoT, one row per run soft-keyed to the item id, capturing ISO-8601-UTC

date_time + tested_by (closed set User|Operator|AI-agent|HelixQA) + result (§11.4.45 PASS|FAIL|SKIP + detail) + in-depth observations (long markdown, facts or explicit UNCONFIRMED: / PENDING_FORENSICS: §11.4.6, the background a future fix needs) + action_taken + status_changed +from/to (did this run change status + WHY) + §11.4.69 evidence_path + feature_class , with a SCHEMA CONSTRAINT making a PASS row WITHOUT a non-empty evidence path impossible (a PASS-bluff rejected by the schema itself); **(b)** BOTH an in-depth diary doc + a derived at-a-glance summary VIEW (total/pass/fail/skip + last-verdict + last-run + status-changes + distinct testers/feature-classes, DERIVED never duplicated §11.4.93) feeding §11.4.132 risk-ordering; **(c)** four-format per-item Diary / Diary_Summary exports §11.4.65 + §11.4.86-drift-proof-fingerprinted (sha256 of the sorted diary keyset) + §11.4.106 docs_chain-bound so a diary row

not exported/pushed FAILs a gate (never-missed); **(d)** external-tracker SUB-TASK model — the item task's description stays the item (§11.4.148 D2, NOT polluted with diary text), each run a CHILD sub-task (type `task`) with a `{TODO|In-progress|Completed}` lifecycle (collapsed per §11.4.33/.112), observations + action + an Evidence: line (path not raw artefact §11.4.10/.13) + a diary-entry idempotency key, reusing the §11.4.148 D5 idempotent rate-limited credential-redacted sink-side-proven push, mapper project-agnostic §11.4.28; **(e)** MINIMAL-LLM deterministic bash/Go tooling (zero LLM in the data path — the `observations` prose is authored by whoever ran the test; tooling only stores/renders/pushes/validates), 100% test-covered §11.4.27 (unit + integration-against-the-real-tracker with an honest §11.4.3 SKIP-when-token-absent never a faked PASS + export + HelixQA Challenge + paired §1.1 mutation) so a diary PASS-bluff is mechanically impossible. Honest boundary §11.4.6 — the diary guarantees a complete auditable per-item test-history, NOT that any single run's verdict is correct (rests on its own §11.4.69/.107/.123 evidence). Composes §11.4.6/.27/.45/.50/.55/.65/.69/.86/.93/.95/.106/.107/.123/.132/.148/.10/.13/.28/.33/.112/ §1.1. Classification: universal (§11.4.17) — consumer supplies its DB path, diary doc layout, export formats, tracker sub-task field map, docs_chain context per §11.4.35. Propagation gate `CM-COVENANT-114-149-PROPAGATION` (literal `11.4.149`) + recommended gates `CM-TEST-DIARY-SYNC` / `CM-DIARY-PASS-REQUIRES-EVIDENCE` + paired §1.1 mutation (gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.149. Non-compliance is a release blocker. No escape hatch — no `--skip-testing-diary`, `--diary-without-evidence`, `--no-diary-summary`, `--diary-without-tracker-subtask`, `--llm-in-diary-data-path`, `--diary-export-optional` flag.

§11.4.150 — Mandatory deep multi-angle web research per change/issue, before declaring fixed or structural (User mandate, 2026-06-11). Verbatim operator mandate: "For every single issue we fix or improvement we make — besides tight systematic-debugging, fixing, code review by independent agents, comprehensive tests — we MUST ALWAYS do deep web research from various angles! No matter how big/small/simple/complex, dig deep the internet for articles, technical documentation, APIs, open-source code — EVERYTHING that can help make the best possible solution OR confirm we don't have a more serious problem we're unaware of! We MUST do everything possible to STOP the constant issue-reopening and finally start closing items as fixed+working! Do this ALWAYS in

parallel with the main work stream!" For EVERY fix / improvement / change / closure — no matter how big, small, simple, or complex — the agent MUST, IN ADDITION to tight systematic-debugging (§11.4.102), the fix, independent-agent code review (§11.4.125 / §11.4.142), and comprehensive multi-layer tests (§11.4.4(b) / §11.4.40), perform a DOCUMENTED **deep multi-angle web research pass** digging the internet from VARIOUS angles — articles, official + vendor technical documentation, API references, standards, issue trackers, maintainer guidance, reusable open-source code — to BOTH (i) discover the best possible solution AND (ii) confirm there is NOT a more serious underlying problem the team is unaware of. Forensic case study (FACT, 2026-06-11): a subtitle-on-secondary-display goal verdicted `Won't-fix: structurally-impossible` (§11.4.112 secure-surface pixel-blanking) was OVERTURNED by a deep multi-angle research pass that surfaced the real OCR-of-the-PRIMARY-display path — a "we can't" became a shipped capability; the canonical anti-pattern of a structural verdict reached for want of research. The mandate (ALL hold): **(A) No closure-as-fixed/structural WITHOUT a documented deep-research pass** — no item may be marked `Fixed / Implemented / Completed` (§11.4.33) OR classified `structurally-impossible won't-fix` (§11.4.112) until a deep multi-angle pass is documented; §11.4.150 makes §11.4.8's pass UNCONDITIONAL (every item, however trivial, at the closure AND structural-verdict gates). **(B) Multiple angles, not a single lookup** — ≥ 2 genuinely-distinct angles (official-docs / standards / known-bug-trackers / alternative-approach / failure-mode / security / performance / platform-constraint / open-source-precedent) so a single-source confirmation-bias miss (§11.4.145) cannot pass; a one-link drive-by is NOT deep multi-angle. **(C) Confirm-no-bigger-problem, not just find-a-fix** — explicitly seek evidence the fix does not mask a deeper defect AND the closure/verdict is genuinely safe; "we found nothing worse" requires the enumerated search (§11.4.118). **(D) Latest-source + cited** — LATEST authoritative versions per §11.4.99 (never training-data/memory), each cited by URL + access date in the research artefact AND the closure commit footer (`Deep-research <date>: <urls>` OR the literal `NO external solution found – original work` per §11.4.8). **(E) Reopen-breaking is the PURPOSE** — STOP the constant Fixed ↔ Reopened churn (§11.4.34 / §11.4.55); a reopen whose root cause a deep pass would have surfaced is a §11.4.150 miss; composes §11.4.7 (demotion-evidence) + §11.4.112 (structural verdict now REQUIRES the cited-authorities pass). **(F) ALWAYS in parallel with the main stream** — background subagent-driven (§11.4.70 / §11.4.20 / §11.4.103 / §11.4.89) concurrent with main fix/build/test work, NEVER serialising it or stalling

the loop (§11.4.94 / §11.4.97 / §11.4.101 / §11.4.126). **(G) Apply ASAP to every workable item** — retroactively + going forward across the §11.4.93 / §11.4.95 SSoT; items closed without a documented pass are re-audited in the §11.4.40 / §11.4.42 release-gate sweep. Honest boundary (§11.4.6): the pass reduces the unknown-unknown surface + breaks the most common reopen causes — it does NOT prove zero remaining defects (§11.4.118) and does NOT replace §11.4.108 runtime-signature verification, §11.4.125 / §11.4.142 review, or §11.4.40 retest; it is the research layer every fix/closure/structural-verdict additionally crosses, and "we probably don't have a bigger problem" without the enumerated multi-angle search is a guess (§11.4.6), never a finding. Composes §11.4.8 / §11.4.99 / §11.4.123 / §11.4.118 / §11.4.145 / §11.4.125 / §11.4.142 / §11.4.7 / §11.4.112 / §11.4.34 / §11.4.55 / §11.4.70 / §11.4.20 / §11.4.89 / §11.4.103 / §11.4.40 / §11.4.42 / §11.4.93 / §11.4.95 / §11.4.6 / §1.1. Classification: universal (§11.4.17) — the consuming project supplies its concrete research corpora, angle set, item tracker, and closure-commit-footer convention per §11.4.35. Propagation gate `CM-COVENANT-114-150-PROPAGATION` (literal `11.4.150`) + recommended gate `CM-DEEP-RESEARCH-PER-ISSUE` (every closed/structural-verdicted item carries a documented multi-angle deep-research artefact + cited-source closure footer, run in parallel, before the closure/structural verdict is accepted) + paired §1.1 mutation (strip the literal → propagation gate FAILs; close an item or reach a `structural` verdict with no documented multi-angle research artefact / cited footer → `CM-DEEP-RESEARCH-PER-ISSUE` FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.150. Non-compliance is a release blocker. No escape hatch — no `--skip-deep-research`, `--single-source-suffices`, `--trivial-change-no-research`, `--close-without-research`, `--structural-without-research`, `--serialise-research`, `--research-from-memory-OK` flag.

§11.4.151 — Project-prefixed release-tag/version-naming mandate (User mandate, 2026-06-12). Verbatim operator mandate: "Every release tag and version name we create — on the main repository and on every Submodule we own — MUST be prefixed with the project's release prefix, e.g. `myproject-1.0.0-dev-0.0.1`. The prefix MUST come from `HELIX_RELEASE_PREFIX` in our `.env` if it is set, otherwise from the lowercased project root directory name. The SAME prefix MUST be used across the main repo and all owned Submodules in one release so a release is greppable across every repository." Every release tag AND every version name created on the main repository AND on every owned-by-us

submodule (§11.4.28) MUST be prefixed with the project's release prefix, form `<PREFIX>-<version>` (canonical example: `myproject-1.0.0-dev-0.0.1`), so a release is identifiable + greppable across every repository it spans (one `git tag -l '<PREFIX>-*'` enumerates the whole release surface). **Prefix resolution order (closed-set, deterministic — §11.4.6 no-guessing):** (1) `HELIX_RELEASE_PREFIX` from the project's `.env` — authoritative when set; `.env` is git-ignored per §11.4.30 and the variable is documented in the tracked `.env.example` (a §11.4.77 re-obtain mechanism, never committed); (2) fallback = the lowercased snake_case form of the project root directory name (no spaces) per §11.4.29 — used whenever the env var is unset/empty, so a prefix is ALWAYS resolvable from the checkout with zero operator input. **The prefix MUST be IDENTICAL across the main repo and all owned submodules within a single release** — a release that tags the main repo `<PREFIX>-1.2.0` while tagging an owned submodule with a different/unprefixed value is a §11.4.151 violation (the cross-repo grep no longer enumerates the release). Version codes increment monotonically within the prefix (`<PREFIX>-...-0.0.1` → `<PREFIX>-...-0.0.2` → ...), never reset, never skipped. Honest boundary (§11.4.6): the prefix guarantees a release is identifiable + uniform across every repository, NOT that its contents are correct — the tag is still created only after the §11.4.40 full-suite retest GREEN and reaches every upstream via the §11.4.113 merge-onto-latest-main path (NEVER a force-push), fanned out per §2.1. Composes §2.1 / §11.4.29 / §11.4.30 / §11.4.40 / §11.4.113 / §11.4.126 (the release-scope terminal condition is a published, prefixed tag). Classification: universal (§11.4.17) — the consuming project supplies its concrete prefix value + the `HELIX_RELEASE_PREFIX` env var per §11.4.35. Propagation gate `CM-COVENANT-114-151-PROPAGATION` (literal `11.4.151`) + recommended gate `CM-RELEASE-PREFIX-NAMING` (every release tag/version on the main repo + every owned submodule carries the resolved `<PREFIX>-` prefix, identical across the release) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; create an unprefixed or differing-prefix release tag → `CM-RELEASE-PREFIX-NAMING` FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.151. Non-compliance is a release blocker. No escape hatch — no `--no-release-prefix`, `--unprefixed-tag`, `--prefix-optional`, `--differing-submodule-prefix` flag.

§11.4.152 — Crashlytics-recorded-data continuous monitoring + systematic-debug + regression-test-coverage mandate (User mandate, 2026-06-13).

Verbatim operator intent: "For every project that has Firebase Crashlytics enabled /

wired, we MUST continuously monitor ALL of the Crashlytics-recorded data — crashes, ANRs, performance traces, and non-fatals — systematically debug each, fix and improve, and cover everything with validation and verification tests. This MUST be checked regularly, with no false results and no bluff of any kind!" Every project that has Firebase Crashlytics enabled/wired (SDK linked into a shipping artifact, crash+non-fatal+ANR reporting active) MUST treat the Crashlytics console as a first-class captured-evidence channel from real end-user devices and continuously process every datum it records — an open Crashlytics issue with a green test suite is a §11.4 PASS-bluff at the field-telemetry layer (a real user hitting a broken feature while the suite reports green). STRENGTHENS+COMPLETES §11.4.47: §11.4.47 owns the periodic REVIEW + dedup + Issue-creation pass that SURFACES Crashlytics/Analytics/Performance findings; §11.4.152 owns what happens to each surfaced item AFTER — the systematic-debug → fix/improve → regression-test-coverage lifecycle that drives it to a proven regression-immune closure (§11.4.47 finds it; §11.4.152 fixes it + proves it stays fixed; both mandatory, neither substitutes). The mandate (ALL hold): (1) **continuous monitoring of ALL four surfaces** — fatal crashes, ANRs, performance traces/regressions, AND non-fatals (skipping any one is a PASS-bluff; non-fatals are the silent class — every catch/fallback/recovered-error path is a real degraded UX and MUST be triaged, not ignored because the app did not crash); (2) **systematic-debugging of each (reproduce-before-fix)** per §11.4.102 Iron Law + §11.4.115 RED-baseline-on-the-broken-artifact (the recorded stacktrace/ANR-trace/non-fatal-context IS the §11.4.5/.69 captured defect-present evidence) BEFORE the fix — a fix without a prior reproducing test is a §11.4.43/.123 violation; unreproducible ⇒ deep-research-before-declaring-untestable §11.4.123/.150; (3) **a fix/improvement for every confirmed issue** against its proven root cause (§11.4.9/.43/.108), "acknowledged in console" is not a fix, muting a recurring issue without a tracked rationale is a §11.4.6/.90 violation; (4) **validation+verification regression-test coverage per closed issue** — in the SAME commit as the fix register a permanent §11.4.135 regression guard (§11.4.115 polarity test: RED_MODE=1 captures the recorded defect on the pre-fix artifact, RED_MODE=0 standing GREEN guard asserting ABSENT) exercising the same user-reachable path the stacktrace identifies, with rock-solid captured evidence §11.4.5/.69/.107/.123 + a closure log citing the console issue id/URL + root-cause + fix commit + the validation/verification test paths; a Crashlytics issue marked resolved WITHOUT a falsifiable regression test is FORBIDDEN (the silent-recurrence vector, §11.4.138 operator-escape class applied to field telemetry); (5) **regular cadence** reusing the §11.4.47 five-trigger set

(pre-build/pre-flash/pre-distribute/pre-tag blocking; daily + post-deployment-burn-in non-blocking) — a one-time sweep never repeated is a violation; (6) **no false results, no bluff** (§11.4.1/6 — "no new issues" requires the enumerated monitored-surfaces+window §11.4.118; "fixed" requires the RED → GREEN flip §11.4.115/.130; a guard whose §1.1 mutation does not FAIL it is a bluff gate). Honest boundary §11.4.6: processing every recorded issue breaks the silent-recurrence vector — it does NOT prove zero remaining field defects nor replace §11.4.108/.40; an issue is "closed" only when its guard is GREEN on a clean artifact, never when the console mark is flipped. Classification: universal (§11.4.17) — the consuming project supplies its Crashlytics handle, console-access credential (§11.4.10, never logged), severity table, and per-issue closure-log path per §11.4.35; the reference project-level instantiation is a consuming project's §6.O (Crashlytics-Resolved Issue Coverage — per-issue validation test + Challenge test + `.lava-ci-evidence/crashlytics-resolved/<date>-<slug>.md` closure log) + §6.AC (Comprehensive Non-Fatal Telemetry — every catch/fallback records a non-fatal with triage context). Composes §11.4.1/.5/.6/.9/.34/.40/.43/.47/.69/.90/.102/.107/.108/.115/.118/.123/.130/.135/.138/.150/ §1.1. Propagation gate `CM-COVENANT-114-152-PROPAGATION` (literal `11.4.152`) + recommended gate `CM-CRASHLYTICS-ISSUE-FULLY-COVERED` (every closed Crashlytics issue carries a systematic-debug root-cause record + a registered §11.4.135 regression guard + a closure-log entry citing the console issue id/URL + the validation/verification test paths; a console-resolved issue with no falsifiable regression guard FAILs) + paired §1.1 meta-test mutation (gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.152. Non-compliance is a release blocker. No escape hatch — no `--skip-crashlytics-monitoring`, `--monitor-crashes-only`, `--skip-non-fatals`, `--resolve-without-regression-test`, `--console-mark-is-fixed`, `--monitor-once`, `--mute-without-rationale` flag.

§11.4.153 — Comprehensive per-feature Status + Status_Summary document set with mandatory video-recording confirmation (User mandate, 2026-06-15).

Every project MUST maintain, under `docs/features/`, a comprehensive feature Status document set (`Status.md` + its §11.4.56 `Status_Summary.md` companion) enumerating EVERY system component, EVERY client app/binary/surface (TUI / CLI / Web / desktop / mobile / API / gRPC / library / submodule / infrastructure), and EVERY feature — including features ported from any incorporated CLI-agent / submodule catalogue (§11.4.74) — with NO single feature

left out. (1) **Total categorized coverage** — per-feature table organized Component → Category, reconciled against the actual codebase (a code-present feature missing from the table, or a row with no code, is a §11.4.6/§11.4.118 bluff).

(2) **Per-feature fields** — Component / Feature / Category / Implementation / Wiring (genuinely end-user-reachable, not merely compiled, §11.4.108) / Real-use / Tests-coverage (four-layer §11.4.4(b)) / Validation (PASS / FAIL / SKIP / PENDING_FORENSICS / OPERATOR-BLOCKED per §11.4.45) / **Video-recording confirmation** (path to the real-use video or honest gap marker). (3) **Mandatory per-feature real-use video** — every user-visible "confirmed" claim backed by a recorded real-use video of a genuine end-user scenario (real prompts → real LLM/ service responses → real results), NEVER a frozen/stale frame (§11.4.107), NEVER faked/mock/demo-loop/bluff response or LLM error passed off as success (§11.4.2/§11.4.5), stored at the project-declared recording path (§11.4.35); a confirmed row with no real video or a bluff video = §11.4 PASS-bluff; autonomous-infeasible ⇒ honest §11.4.3/§11.4.52 SKIP + tracked migration item, NEVER a faked confirmation. (4) **Video-analysis remediation loop** — every video analysed; any defect it surfaces triggers §11.4.102 systematic-debug → fix → §11.4.146 retest → re-record → clean GO per §11.4.134 before "confirmed" (a video exposing a broken feature is a §11.4.4 test-interrupt). (5) **Always-in-sync** — §11.4.45-class roster/corpus-backed, §11.4.106 docs_chain-bound + §11.4.86 drift-proof fingerprint (sha256 of sorted feature-key roster AND sorted video-artefact roster, NOT mtime), re-syncs out-of-the-box; stale = violation. (6) **Four-format export** — HTML + PDF + **DOCX** (this doc class ADDS DOCX to the §11.4.65 HTML+PDF set; other classes unchanged), in sync per §11.4.60. (7) Follows §11.4.44/.45/.56/.57/.59/.60. (8) **MP4 format REQUIRED**. All video confirmations MUST be in .mp4 format (H.264). Window-specific capture ONLY (§11.4.159(A)). Vision validation REQUIRED (§11.4.159(D)). .cast files are supplementary only. Honest boundary §11.4.6 — guarantees a complete video-confirmed always-synced ledger, NOT per-feature correctness (rests on each row's §11.4.69/.107/.123 evidence) and NOT a §11.4.40 retest substitute. Classification: universal (§11.4.17). Composes §11.4.2/.5/.44/.45/.52/.56/.57/.59/.60/.65/.86/.102/.106/.107/.108/.118/.123/.134/.146/ §1.1. Propagation gate CM-COVENANT-114-153-PROPAGATION (literal 11.4.153) + recommended gates CM-FEATURE-STATUS-COMPLETE + CM-FEATURE-STATUS-VIDEO-CONFIRMED + paired §1.1 mutation.

Canonical authority: constitution submodule `Constitution.md` §11.4.153. Non-compliance is a release blocker. No escape hatch — no `--skip-feature-status`, `--feature-without-video`, `--frozen-video-OK`, `--bluff-response-video-OK`, `--skip-video-analysis`, `--feature-ledger-incomplete-OK`, `--no-docx-export`, `--allow-feature-ledger-drift` flag.

§11.4.154 — Window-scoped capture + fresh-corpus rotation for feature/QA recordings (User mandate, 2026-06-15). Verbatim: "recording you perform does record the window containing apps and services, not the whole desktop or monitor screen!" + "All old recording files MUST BE removed when new one starts!" Refines §11.4.2/.5/.107/.153 recording discipline with two capture-hygiene invariants. **(A) Window-scoped, NOT whole-screen** — every feature/QA video MUST capture ONLY the window/surface of the app/service under test (GUI window / CLI-TUI terminal pane / web tab-viewport / device-emulator-simulator frame), NEVER the whole desktop/monitor or unrelated windows; whole-desktop capture leaks operator-private content (§11.4.10/.83), dilutes the §11.4.107 liveness/freeze oracle, and breaks the §11.4.137 OCR/ROI content oracle. Target the window/region by stable identity (window id/title, device serial, browser context, tmux target) per §11.4.111 — never a fixed full-screen device index. Platform genuinely cannot capture below whole-screen ⇒ honest §11.4.3 SKIP + tracked migration item, never a whole-screen pass-off. **(B) Fresh-corpus rotation** — when a new recording run for a scope begins, the agent's OWN prior in-scope stale recordings at the raw recording path MUST be removed FIRST so the live corpus reflects the current run (§11.4.107 not-stale + §11.4.86 roster-freshness). Honest boundary (§11.4.6 + §9.2): "remove old" = the agent's own prior recordings for the SAME scope/project ONLY — NEVER another project's/scope's/operator-authored files; uncertain ⇒ surface, don't delete (§11.4.122); committed `docs/qa/<run-id>/` evidence (§11.4.83) is the durable record, NOT rotated. Classification: universal (§11.4.17). Composes §11.4.2/.5/.10/.83/.86/.107/.111/.122/.128/.137/.153/§9.2/.6. Propagation gate `CM-COVENANT-114-154-PROPAGATION` (literal `11.4.154`) + recommended gate `CM-WINDOW-SCOPED-FRESH-CORPUS-RECORDING` + paired §1.1 mutation.

Canonical authority: constitution submodule `Constitution.md` §11.4.154. Non-compliance is a release blocker. No escape hatch — no `--whole-screen-capture-OK`, `--skip-window-scope`, `--keep-stale-recordings`, `--no-corpus-rotation`, `--full-desktop-recording` flag. **(C) MP4 auto-conversion**

REQUIRED. Any `.cast` file produced MUST be auto-converted to `.mp4` via `agg + ffmpeg` immediately after capture. The `.mp4` is the primary evidence; `.cast` is supplementary only.

§11.4.155 — Project-name-prefixed feature/QA recording filenames (User mandate, 2026-06-15). Verbatim: "All recorded videos MUST START with prefix: the PROJECT NAME (ALWAYS USE THE PROJECT NAME). Project name MUST be obtained according to the constitution's own project-name resolution." Every recorded video the project produces — every feature/QA real-use recording (§11.4.153), every window-scoped capture (§11.4.154), every always-on device recording (§11.4.128), every raw or curated artefact at the project-declared recording path (§11.4.35) + the committed `docs/qa/<run-id>/` trail (§11.4.83) — MUST have a filename that STARTS WITH the PROJECT-NAME prefix, ALWAYS; an unprefixed recording is a §11.4.155 violation (a multi-project/scope corpus on one host per §11.4.128/103 becomes un-greppable + un-attributable — the §11.4.151 identify-and-grep failure on the recording-corpus axis). **Prefix resolution (closed-set, deterministic — §11.4.6, IDENTICAL to §11.4.151):** (1)

`HELIX_RELEASE_PREFIX` from `.env` (authoritative, git-ignored §11.4.30, documented in tracked `.env.example` §11.4.77) else (2) lowercased snake_case project-root dir name §11.4.29 — ALWAYS resolvable, zero operator input. SAME prefix for EVERY recording in a checkout so `ls '<PREFIX>--- '*` enumerates the corpus; canonical form `<PREFIX>---<feature-or-scope>---<run-id>.<ext>` (`---` delimits the prefix unambiguously); MUST equal the §11.4.151-resolved release-tag prefix (divergence is itself a §11.4.155 violation — one project, one name). Honest boundary (§11.4.6): the prefix guarantees attribution + greppability, NOT content validity (still §11.4.107/137/153) and does NOT relax §11.4.154's window-scope/rotation (rotation removes the agent's OWN `<PREFIX>--- *` only; foreign/operator files surfaced not deleted §11.4.122/§9.2). Classification: universal (§11.4.17). Composes §11.4.151/128/153/154/111/83/6/29/30/35/77/86/§1.1. Propagation gate `CM-COVENANT-114-155-PROPAGATION` (literal `11.4.155`) + recommended gate `CM-RECORDING-PROJECT-NAME-PREFIX` + paired §1.1 mutation.

Canonical authority: constitution submodule `Constitution.md` §11.4.155. Non-compliance is a release blocker. No escape hatch — no `--no-recording-prefix`, `--recording-without-project-name`, `--unprefixed-recording`, `--prefix-optional-for-recording`, `--differing-recording-prefix` flag.

§11.4.156 — All CI/CD automation (GitHub Actions / GitLab pipelines / equivalents) MUST be disabled (User mandate, 2026-06-15). Verbatim operator mandate: "Any GitHub actions or GitLab pipelines MUST BE disabled! Add this critical mandatory rule / mandatory constraint into the root constitution Submodule, commit and fetch all its changes to all upstreams and make sure we respect and follow this rule (we do apply it) ASAP!!!" Every repository this Constitution governs — main repo, this constitution submodule, every owned + nested submodule we author and push — MUST ship with ALL server-side CI/CD automation DISABLED: no push to any owned upstream may trigger a GitHub Actions run, GitLab pipeline, or equivalent (Jenkins/CircleCI/Travis/Drone/Woodpecker/Bitbucket/Azure, any on: push / schedule / workflow_dispatch). GENERALISES + makes ABSOLUTE the §11.4.75 Layer-5 posture (remote CI DISABLED, workflow preserved at a ...disabled-local-only non- .yaml name a provider ignores) across ALL governed repos; enforcement migrates to the LOCAL §11.4.75 git-hook ritual + §11.4.40 pre-tag sweep, never a remote runner. ALL hold: **(A)** zero active .github/workflows/*.yaml|yml / .gitlab-ci.yml / .gitlab/** / equivalent at the ROOT of any governed repo/submodule (the only place a provider executes); **(B)** "disabled" = a push triggers ZERO runs — delete OR rename to a non-trigger name (§11.4.75 .disabled / .disabled-local-only); a live- on: + if:false workflow still queues runs, NOT compliant; **(C)** scope = repos we author+push — vendored/third-party nested configs below the root (AOSP external/** , prebuilts/** , vendored submodules) are INERT (a provider never runs a non-root config), OUT of scope, MUST NOT be mass-edited (§11.4.29 vendor-exempt); test = "does a push to OUR upstream trigger a run?" yes⇒disable, inert⇒document+leave (§11.4.6 verify-not-assume); **(D)** no new CI may be added (release blocker); **(E)** pre-push verify git ls-files | grep -E '^\.github/workflows/.*\.ya?ml\$|^\.gitlab-ci\.yaml\$' empty for authored repos, §11.4.109-class PreToolUse guard + gate enforce mechanically. Honest boundary (§11.4.6): file-level disabling stops FILE-triggered runs, NOT provider-side server settings (org-default required workflows, branch-protection required checks, provider scheduled exports) — the operator turns those off; the agent documents what it cannot reach, never claims unachieved completeness. Composes §11.4.75/.29/.6/.40/.42/.109/.113/§2.1. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-156-PROPAGATION (literal 11.4.156) + recommended gate CM-NO-ACTIVE-CI + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; add a root .github/workflows/x.yaml to an

authored repo → `CM-NO-ACTIVE-CI FAILS`). **Canonical authority:** constitution submodule `Constitution.md` §11.4.156. Non-compliance is a release blocker. No escape hatch — no `--allow-ci` , `--enable-workflow` , `--keep-pipeline` , `--remote-ci-OK` , `--ci-exempt` flag.

§11.4.157 — GEMINI.md maintained in lockstep with CLAUDE.md /

AGENTS.md / QWEN.md (User mandate, 2026-06-15).

Verbatim operator mandate: "Make sure with CLAUDE.md, AGENTS.md, QWEN.md we maintain GEMINI.md too! Add this mandatory fact / rule to root constitution Submodule we are inheriting / extending - CONSTITUTION.md, CLAUDE.md, QWEN.md, AGENTS.md, GEMINI.md and other related relevant files!" Forensic FACT (2026-06-15): when §11.4.156/§11.4.157 were authored, Constitution/CLAUDE/AGENTS/QWEN carried the family through §11.4.155 but GEMINI.md had silently drifted to §11.4.141 — 14 mandates (§11.4.142–155) never propagated — a §11.4 propagation-bluff (a Gemini-CLI agent reads a stale Constitution). GEMINI.md is a FIRST-CLASS governance context carrier EQUAL to CLAUDE.md/AGENTS.md/QWEN.md, never optional/best-effort. ALL hold: **(A)** five-carrier lockstep — no governance change is complete until GEMINI.md carries it alongside the other three mirrors; GEMINI.md is added to the §11.4.26 propagation + cross-reference set explicitly; **(B)** no silent drift — GEMINI.md lagging the other mirrors' highest rule is a §11.4.157 violation (§11.4.65-class), back-fill required; **(C)** equal status — GEMINI.md restates the SAME literal `11.4.N` anchors the propagation gates require (§11.4.35), fleet count INCLUDES GEMINI.md; **(D)** consumer projects' own CLAUDE/AGENTS/QWEN/GEMINI bind too (§11.4.35). Honest boundary (§11.4.6): the §11.4.142–155 GEMINI.md back-fill is a tracked release-blocking remediation; claiming GEMINI.md "in sync" while the back-fill is incomplete is itself a §11.4.157 violation. Composes §11.4.26/.35/.17/.44/.65/.140/.156/§1.1. Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-157-PROPAGATION` (literal `11.4.157` , GEMINI.md INCLUDED) + recommended gate `CM-GEMINI-MD-LOCKSTEP` + paired §1.1 meta-test mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.157. Non-compliance is a release blocker. No escape hatch — no `--skip-gemini-md` , `--gemini-optional` , `--gemini-lag-OK` , `--four-carrier-suffices` flag.

§11.4.158 — Intensive all-feature/flow/edge-case video-recording + read-the-screen content-verification mandate (User mandate, 2026-06-16). Every project MUST be covered by intensive automated testing that exercises + RECORDS every feature/flow/use-case/edge-case (valid/invalid/boundary/concurrent/failure-

injection §11.4.85), each recording showing the feature GENUINELY WORKING with REAL results (real prompts → real responses → real outputs §11.4.153) and NO false/simulated/stale/frozen result (§11.4.2/.5/.107) — a recording showing an error/bluff/non-working feature is a FINDING (→ §11.4.153(4)/§11.4.4 fix → retest → re-record), never a confirmation. The testing System MUST ACTUALLY READ every shown log line / message / UI label / dialog / toast / status text and VERIFY it is a genuine working result (OCR/ROI §11.4.117/.137 + confidence floor for pixel surfaces; direct capture for terminal/log; §11.4.107(10) self-validated golden-good/golden-bad analyzer) — "a video was produced" is NOT evidence, "the System read the screen + confirmed a genuine result" is. HelixQA (§11.4.27) MUST drive this exercise → record → read → score pass, PASS only on a read-confirmed genuine result + captured artefact path (§11.4.69). Default recording save path = `$HOME/Downloads` (host user's home Downloads, resolved at runtime never hardcoded) unless a project declares an override per §11.4.35; §11.4.155 project-prefix + §11.4.154 window-scope/rotation + §11.4.128 git-ignored raw corpus + §11.4.83 curated docs/qa evidence apply. **Vision analysis MANDATORY for every recording** — after every recording, the agent MUST read the terminal output / video content and verify the feature actually works. A recording without a vision-confirmed verdict is a §11.4.158 violation. Composes

§11.4.2/.5/.25/.27/.52/.69/.83/.85/.107/.108/.117/.118/.128/.137/.138/.153/.154/.155/
§1.1. Classification: universal (§11.4.17). Propagation gate

`CM-COVENANT-114-158-PROPAGATION` (literal `11.4.158`) + recommended gates

`CM-INTENSIVE-RECORDING-COVERAGE` +

`CM-RECORDING-CONTENT-READ-VERIFIED` + paired §1.1 mutation. **Canonical**

authority: constitution submodule `Constitution.md` §11.4.158. Non-compliance is a release blocker. No escape hatch — no `--skip-recording-coverage`, `--video-without-content-read`, `--happy-path-recording-suffices`, `--recording-path-anywhere`, `--unread-recording-OK`, `--skip-helixqa-read` flag.

§11.4.159 — Mandatory window-specific video recording + vision validation

mandate (User mandate, 2026-06-20). Every feature test, validation, verification, challenge, and QA session that produces video evidence MUST comply with ALL of the following: **(A) Window-specific recording ONLY** — every video MUST record ONLY the target application window (Terminal pane, TUI app, browser tab, emulator frame), NEVER the whole desktop/monitor screen; use macOS `screencapture -l<window_id>`, Linux `xdotool + ffmpeg`, or equivalent; whole-desktop capture

leaks operator-private content (§11.4.10), dilutes the §11.4.107 liveness oracle, and breaks §11.4.137 OCR/ROI; platform genuinely cannot capture below whole-screen => honest §11.4.3 SKIP + tracked migration item. **(B) MP4 format REQUIRED** — `.mp4` (H.264, `movflags +faststart`, `pix_fmt yuv420p`); `.cast` files are supplementary only; auto-convert via `agg + ffmpeg` per §11.4.154(C). **(C) Project-name prefix REQUIRED** — filename MUST start with project name in snake_case from `HELIX_RELEASE_PREFIX` in `.env` (§11.4.151) or lowercased project root dir name (§11.4.29); format `<project_name>-<feature>-<timestamp>.mp4`; unprefixed = violation. **(D) Mandatory vision validation** — after EVERY recording, the agent MUST read the terminal output / video content and verify: (i) feature ACTUALLY WORKS, (ii) LLM responses are REAL, (iii) all tests show PASS, (iv) no "TODO implement" / "simulate" / "for now" patterns, (v) output demonstrates end-user working feature; MUST produce a verdict (PASS/FAIL) with evidence path (§11.4.69). **(E) Terminal window cleanup** — after each recording, dismiss/close ONLY the Terminal window used for that recording (window-specific close via `osascript` / `xdotool`); MUST NOT close windows belonging to other processes. **(F) Real results ONLY** — every recording MUST show REAL working features; errors/empty output/simulated responses trigger fix → retest → re-record. **(G) Re-runnable evidence** — the command shown MUST be re-runnable to produce the same results. **(H) Fresh-corpus rotation** — remove agent's own prior in-scope stale recordings FIRST (§11.4.154). **(I) Content verification MANDATORY — not duration-based** — the value of a recording is NOT its duration but its CONTENT; a recording MUST demonstrate the ACTUAL feature being used with REAL results; before accepting ANY recording, the agent MUST verify: expected output patterns ARE present (e.g., test PASS lines, API response data, feature-specific output), the feature ACTUALLY WORKS as demonstrated (not just "something ran"), LLM responses are REAL content (not simulated, not placeholder, not empty), every claim of "working" is backed by visible evidence; a 5-second recording showing a feature working correctly is MORE valuable than a 60-second recording of empty terminal; duration is NOT a proxy for quality. **(J) Expected-content specification REQUIRED** — before recording, the agent MUST specify what content SHOULD appear (expected patterns, expected test results, expected API responses); after recording, the agent MUST verify these patterns ARE present; if expected content is MISSING, the recording is REJECTED regardless of duration. **(K) Content-verification recording workflow** — mandated workflow: (1) SPECIFY expected content patterns, (2) RECORD the feature execution, (3) EXTRACT all text from the recording, (4) VERIFY expected patterns

are present, (5) CHECK for simulated/placeholder content, (6) ACCEPT only if ALL patterns found AND zero bluffs detected, (7) REJECT and re-record if ANY pattern missing or bluff detected. **(L) Root cause analysis REQUIRED for rejected recordings** — when a recording is rejected (missing expected content, bluff detected, or empty capture), the agent MUST investigate WHY before re-recording per §11.4.102; determine the root cause (timing issue, wrong command, tool failure, etc.) and fix it; simply re-recording without understanding WHY the first attempt failed is a §11.4.159 violation. **(M) Real-time monitoring RECOMMENDED** — for complex features, use real-time monitoring that analyzes output DURING recording (not after); this catches issues immediately and allows corrective action before the recording completes. Classification: universal (§11.4.17). Composes §11.4.2/.3/.5/.10/.29/.69/.83/.107/.111/.128/.137/.151/.153/.154/.155/.158/§1.1. Propagation gate CM-COVENANT-114-159-PROPAGATION (literal 11.4.159) + recommended gate CM-WINDOW-VIDEO-VALIDATED + paired §1.1 mutation. **Canonical authority:** constitution submodule Constitution.md §11.4.159. Non-compliance is a release blocker. No escape hatch — no --whole-screen-ok , --cast-only , --skip-vision-validation , --no-cleanup , --simulated-recording-ok , --unprefixed-recording flag.

§11.4.160 — Vision-verified recording + HelixQA bridge mandate (User mandate, 2026-06-21). Compact summary: every video recording for feature/QA evidence MUST be processed through a vision/OCR pipeline that reads on-screen content and confirms expected results BEFORE acceptance; the recording system MUST provide a bridge feeding captured frames to HelixQA's test infrastructure (or equivalent) for automated read-the-screen verification against SPECIFY-phase expected patterns (§11.4.159(J)); capture frames at ≤5s intervals, run OCR/vision with self-validated analyzer (§11.4.107(10)), compare text against patterns, produce per-frame PASS/FAIL with evidence path, surface failures immediately for re-record (§11.4.159(L)). Project-configured parameters per §11.4.35. Honest boundary: does NOT replace §11.4.5 quality analysis nor §11.4.108 runtime-signature; FLAG_SECURE uses §11.4.117 proxy oracle. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-160-PROPAGATION + recommended gate CM-VISION-VERIFIED-RECORDING-BRIDGE + paired §1.1 mutation.

§11.4.161 — Rootless container runtime mandate (User mandate, 2026-06-21). Compact summary: every project MUST use Podman rootless (or equivalent) for ALL containerized workloads — rootful Docker/sudo FORBIDDEN unless §11.4.112 platform constraint documented; vasic-digital/containers Submodule

(§11.4.76) sole orchestration layer via `pkg/boot / pkg/compose / pkg/health` ; extend upstream per §11.4.74 for missing capabilities; integration tests boot infra on-demand. Honest boundary: does NOT replace §11.4.10 credentials nor §12.3 hygiene. Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-161-PROPAGATION` + recommended gate `CM-ROOTLESS-CONTAINER-RUNTIME` + paired §1.1 mutation.

§11.4.162 — OpenDesign UI design system mandate (User mandate, 2026-06-21). Compact summary: every project with user-facing interfaces MUST use OpenDesign (<https://github.com/nexu-io/open-design>) as mandatory UI design system — NOT ad-hoc CSS; install as dependency, use token system for all color (light+dark from brand assets), typography, spacing, component tokens; extend upstream (§11.4.74) for unsupported patterns; every UI component ships light+dark + token-diff regression tests via visual regression; no element overlap/font collision/label overlay. Honest boundary: does NOT replace §11.4.27 functional testing nor §11.4.48/49 UI testing. Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-162-PROPAGATION` + recommended gate `CM-OPENDESIGN-UI-SYSTEM` + paired §1.1 mutation.

§11.4.163 — Universal Media Validation (User mandate). Every recorded artifact MUST pass MEDIA VALIDATION pipeline — OCR/transcription/text parsing, compare vs SPECIFY-phase patterns (§11.4.159(J)), self-validated analyzer (§11.4.107(10)), structured verdict, pinpoint on FAIL, REAL-TIME trigger option. Paired §1.1 mutation on golden-bad fixture. Classification: universal. Propagation gate `CM-COVENANT-114-163-PROPAGATION` + paired §1.1 mutation.

§11.4.164 — Universal Auto-Propagation Hook (User mandate). Every constitutional pull triggers `post_update_hook.sh` — detects changed files, registers skills/MCP/hooks/scripts, logs summary. Consumer invokes after every pull. Classification: universal. Propagation gate `CM-COVENANT-114-164-PROPAGATION` + paired §1.1 mutation.

§11.4.165 — Universal Independent Verification Agent (User mandate). Every code/media/docs/config output passes independent verifier (§11.4.70/.20), zero-finding GO (§11.4.134). CODE: review+build+mutations+runtime-signature. MEDIA: pipeline+genuine+format. DOCS: exports+revision. CONFIG: syntax+schema+leaks. Self-validated §1.1 mutation. Classification: universal. Propagation gate `CM-COVENANT-114-165-PROPAGATION` + paired §1.1 mutation.

§11.4.166 — REPEALED (operator decision, 2026-06-22). The Universal Semgrep static-analysis mandate is repealed — Semgrep is NO LONGER mandatory. Static analysis remains encouraged, not mandated. Scaffolding, `submodules/semgrep`, `MCP/pre-commit/PATH` wiring, `docs_chain` `semgrep_status` context, and the `CM-COVENANT-114-166-PROPAGATION` / `CM-SEMGREP-WIRED` gates removed. Anchor 11.4.166 retired. See constitution `Constitution.md` §11.4.166.

§11.4.167 — Big-work-item feature work-stream lifecycle (User mandate, 2026-06-23). Every BIG work item (new feature OR drastic/large fix) MUST be its own isolated feature work-stream — own sibling project copy, own `feature/<slug>` branch, own per-feature flashable builds + feature tags, SEPARATE from trunk until operator manually approves after testing, trunk merged INTO every stream regularly. (A) one big item = one sibling project copy (`<project>_<slug>`); small changes stay on trunk. (B) disk-feasible CoW/reflink clone MANDATORY (`cp -a --reflink=auto` on btrfs/XFS/ZFS/APFS, else `git worktree`); naive deep copies FORBIDDEN; per-stream `out/` + shared ccache; ~0-disk verified (§11.4.69) before relied on (§11.4.6). (C) own branch + greppable tags (§11.4.151 `<base>-feat-<slug>`), NEVER merged to trunk until §11.4.40 GREEN on clean target → §11.4.41 merge-first → §11.4.113 ff-only. (D) trunk merged INTO every live stream per trunk-tag/daily (merge, never rebase). (E) feature branch + tag cascades to EVERY touched submodule (§11.4.113). (F) single-builder build QUEUE (§11.4.58/.12.7/.12.8) + finite-device queue (§11.4.119); idle `out/` GC'd (§11.4.77). (G) every change crosses full gauntlet (§11.4.142/.125/.134/.145 + anti-bluff). (H) stream registry (§11.4.116/.147) + atomic `.partial` → rename + §11.4.84 resume + §9.2 retire-backup; bounded §12/.12.6. Honest boundary (§11.4.6): isolation+disk-feasibility+greppable-identity+merge-discipline, NOT correctness (still §11.4.108/.40). Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-167-PROPAGATION` (literal `11.4.167`) + gates `CM-FEATURE-WORKSTREAM-COW-CLONE` / `-NO-MERGE-UNTIL-APPROVED` / `-TRUNK-SYNC-CADENCE` / `-SUBMODULE-CASCADE` + paired §1.1 mutations.

Canonical authority: constitution `Constitution.md` §11.4.167.

§11.4.168 — Exported-document independent content + textual + full-visual validation mandate (User mandate, 2026-06-23). Every generated/exported document (HTML/PDF/DOCX/any §11.4.65 export-scope format) MUST pass INDEPENDENT validation by a reviewer structurally separate from the generator (§11.4.70/.142, NEVER author self-check), on BOTH source AND every exported

artifact, ALWAYS, across THREE layers: (1) CONTENT — export faithfully carries source intent+data, nothing dropped/truncated/garbled; (2) TEXTUAL — human-readable, NO raw markup/diagram-source (Mermaid gantt / graph / flowchart / sequenceDiagram etc.)/unrendered code-fence leaking as body text (the exact 2026-06-23 raw-gantt-in-PDF failure); (3) FULL VISUAL — diagrams render as IMAGES not source, layout intact, no overlap/cut-off/garble, verified by RENDERING the export (pdftotext catches raw source + pdfimages confirms rendered images + pdftoppm → OCR confirms human-readable content) with captured evidence §11.4.5/.69/.107. Anti-bluff: rubber-stamp ≠ validation; analyzer self-validated golden-good/golden-bad §11.4.107(10) (golden-bad seeded with raw gantt MUST FAIL); iterate to clean GO §11.4.134. Forensic FACT (2026-06-23): a "green" Mermaid fix shipped raw gantt source as readable text into user PDFs because validation grepped HTML for class="mermaid" and NEVER checked the exported PDF — operator caught it (§11.4.138). Honest boundary (§11.4.6): confirms the exported artifact is faithful/readable/visually-rendered — does NOT prove source content is correct (§11.4.142/.135) nor replace the §11.4.65 mtime freshness gate (this is the READABILITY/FIDELITY layer the mtime gate cannot see); FLAG_SECURE/un-rasterisable surfaces document the gap §11.4.112 + §11.4.117 proxy, never a faked PASS. Classification: universal (§11.4.17). Composes §11.4.65/.73/.107/.117/.135/.138/.142/.159/.163/.165/§1.1. Propagation gate CM-COVENANT-114-168-PROPAGATION (literal 11.4.168) + recommended gate CM-EXPORTED-DOC-VISUALLY-VALIDATED + paired §1.1 mutation. **Canonical authority:** constitution Constitution.md §11.4.168.

§11.4.169 — Mandatory comprehensive test-type coverage with anti-bluff captured evidence (User mandate, 2026-06-25). Every project MUST cover a CLOSED enumerated test-type set, each to as-close-to-100% as the domain permits, every PASS citing rock-solid captured PHYSICAL evidence (§11.4.5/.69/.107), zero false results/zero bluff (§11.4/.1/.6): (1) unit (only layer mocks allowed §11.4.27(A)); (2) integration (real System, infra via containers submodule §11.4.76); (3) e2e; (4) full-automation (autonomous, re-runnable, no manual step §11.4.25/.52/.98, deterministic §11.4.50); (5) Challenges (challenges submodule banks §11.4.27(B)); (6) HelixQA (helix_qa submodule — written test banks/suites per app/service/platform + comprehensive fully-autonomous QA sessions §11.4.27); (7) DDoS/load-flood (clean refuse or graceful degrade); (8) security (authz/injection/secret-leak §11.4.10/crypto/CVE/default-deny + security submodule); (9) stress+chaos (load+contention+boundary+failure-injection,

categorised recovery §11.4.85); (10) concurrency/atomicity (no lost updates, transactional atomicity, idempotent replay); (11) race-condition/deadlock (race detector + lock-order analysis; no blocking op in shared-lock region §1); (12) memory (leak census + peak-RSS ceiling + 24h/N-iter soak, no unbounded growth); (13) benchmarking/performance (p50/p95/p99 + throughput vs recorded baseline). The ONLY permitted absence is an honest §11.4.3 SKIP-with-reason, never a silent gap; four-layer §11.4.4(b) enforcement + coverage ledger (§11.4.25/.52); missing required type or OPERATOR_ATTENDED_ONLY = release blocker. Strict enumerated expansion of §11.4.27. Composes §11.4.4/.5/.6/.25/.27/.50/.52/.69/.76/.85/.98/.107/.123/.135/.146/§1.1. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-169-PROPAGATION (literal 11.4.169) + recommended gate CM-MANDATORY-TEST-TYPES-COVERED + paired §1.1 mutation. **Canonical authority:** constitution Constitution.md §11.4.169.

§11.4.170 — Device-independent host-side rendered-UI visual-proof mandate (User mandate, 2026-06-25). Compact summary: every change to ANY user-facing UI surface (screen/component/view/widget/layout/theme on any platform — Android-Compose, iOS-SwiftUI/UIKit, web, desktop, TUI) MUST be proven by DEVICE-INDEPENDENT host-side RENDERED PIXELS before it may be claimed correct — the real component rendered to a PNG ON THE HOST (NO device, NO emulator, NO running app) via a host-render harness (Compose → Roborazzi/ Paparazzi on the JVM; web → Playwright/Storybook snapshot; SwiftUI → snapshot-testing; equivalent per stack), for EVERY relevant screen×state×{light,dark} theme, with the PNG validated by BOTH (i) golden image-diff (per-pixel/perceptual PASS/ FAIL) AND (ii) an OCR/vision oracle reading rendered text+labels+control bounds to assert legibility + NO overlap / NO label-over-label / NO clipping / NO off-screen control / NO collapsed-or-giant-unbounded widget (the §11.4.162 layout invariants PROVEN on real pixels). Forensic FACT (2026-06-25): UI work was "validated" by VALUE-EQUALITY unit tests (a hex colour string / an sp/dp size integer / a design-token field == a constant) that NEVER RENDERED A PIXEL — GREEN while the operator opened a broken unbounded giant-button screen. A value/token-equality / property-assertion UI test is FORBIDDEN as the PROOF a UI is correct (it may SUPPLEMENT, NEVER SUBSTITUTE the rendered-pixel proof); a "green" UI suite with no rendered-pixel artefact for the changed surface is a §11.4 PASS-bluff at the visual layer. Self-validated golden-good/golden-bad analyzer §11.4.107(10), thresholds calibrated on the project's own fixtures §11.4.6. "device offline" is

NEVER a valid skip — host-render IS the device-independent path; honest §11.4.3 SKIP only where the platform genuinely has no host-render harness (tracked migration item, never a value-only fallback). Honest boundary §11.4.6: host-render proves the COMPONENT renders correctly device-independently per commit — it does NOT prove the running app on real hardware (that remains §11.4.153/.158/.159/.160's live-device recording + read-the-screen layer, which §11.4.170 COMPLEMENTS not replaces — both mandatory), does NOT replace §11.4.162 OpenDesign token/theme governance (§11.4.170 PROVES the rendered RESULT of those tokens), DISTINCT from §11.4.168 exported-document visual validation. Classification: universal (§11.4.17). Composes §11.4.3/.5/.27/.40/.69/.107/.108/.117/.153/.158/.159/.160/.162/.163/.168/.169/§1.1. Propagation gate `CM-COVENANT-114-170-PROPAGATION` (literal `11.4.170`) + recommended gate `CM-HOST-RENDERED-UI-VISUAL-PROOF` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.170. Non-compliance is a release blocker. No escape hatch — no `--value-assertion-proves-ui`, `--skip-rendered-pixel-proof`, `--token-equality-suffices`, `--device-offline-skip-visual`, `--single-theme-only`, `--unvalidated-image-diff-OK`, `--no-ocr-layout-oracle` flag.

When in doubt

- Read `constitution/Constitution.md` for the canonical text of every rule.
- Cross-reference the sibling carriers (`CLAUDE.md` / `AGENTS.md` / `QWEN.md`) — they carry the same universal content for their respective agents.
- If still unclear, ask the operator. Do NOT guess. Do NOT bluff (§11.4.6).

§11.4.171 — Mandatory comprehensive human-readable workable-item descriptions (User mandate, 2026-06-29). Every workable item (ATM-NNN ticket) in the project's workable-items database, `Issues.md`, `Fixed.md`, `Issues_Summary.md`, `Fixed_Summary.md`, and ALL derived documents MUST carry a comprehensive human-readable description that explains the item in plain, non-technical language suitable for non-developers (project managers, stakeholders, team leads, business partners). The description MUST be 5-7 sentences minimum, written so that ANY person — regardless of technical background — can understand: (a) WHAT the item is (the feature, fix, or task in simple terms); (b) WHY it matters (the user/business value); (c) HOW it works (the

mechanism, without code references); (d) WHO benefits (the end user or stakeholder); (e) WHAT the expected outcome is (the measurable result). The description is MANDATORY — an item without a human-readable description is a §11.4 violation at the documentation layer. Descriptions MUST NOT use jargon, acronyms without expansion, code references, or technical implementation details as the primary explanation. Technical details MAY appear as supplementary context AFTER the plain-language explanation. The `description` column in the SQLite database (`docs/workable_items.db`) is the SINGLE SOURCE OF TRUTH for item descriptions; all derived documents (Issues.md, Fixed.md, summaries, exports) MUST carry the SAME description text. Pre-build gate `CM-HUMAN-READABLE-DESCRIPTION` (when implemented) walks every item in the DB and asserts the `description` field exists, is non-empty, and is ≥ 50 characters. Paired §1.1 meta-test mutation strips a description → gate FAILs. Classification: universal (§11.4.17). Composes §11.4.15 (item-status tracking), §11.4.16 (item-type tracking), §11.4.19 (column alignment), §11.4.91 (summary clarity), §11.4.93 (SQLite SSoT), §11.4.148 (item integrity), §11.4.65 (universal export). Propagation gate `CM-COVENANT-114-171-PROPAGATION` (literal `11.4.171`) + paired §1.1 mutation.

§11.4.172 — Mandatory production-readiness planning with realistic timeline projection (User mandate, 2026-06-29). Every project under this Constitution MUST maintain a living production-readiness planning document that: (a) provides REALISTIC timeline projections for all product milestones (MVP, beta, production-ready) based on actual development velocity data (items completed per week, reopening rate, blocking dependencies); (b) identifies ALL danger zones, weaknesses, and risk areas with mitigation strategies; (c) accounts for HW delivery timelines, licensing delays, and external dependency risks; (d) defines the critical path and identifies which items can slip without affecting MVP; (e) includes workstation/build-capacity analysis and optimization recommendations; (f) is updated at least monthly or whenever the item count changes by $\geq 10\%$; (g) is cross-referenced against the workable-items database for accuracy. The planning document MUST live at `docs/research/production_planning_<YYYYMMDD>/ANALYSIS.md` with HTML+PDF exports per §11.4.65. Honest boundary (§11.4.6): projections are estimates based on measured velocity — they guarantee the methodology is sound, NOT that specific dates will be met. Classification: universal (§11.4.17). Composes §11.4.6 (no-guessing — projections based on data, not

wishful thinking), §11.4.40 (full retest before tag), §11.4.42 (iteration discipline), §11.4.93 (SQLite SSoT), §11.4.108 (four-layer verification). Propagation gate `CM-COVENANT-114-172-PROPAGATION` (literal `11.4.172`) + paired §1.1 mutation.

§11.4.173 — Containerized + distributed build mandate (User mandate, 2026-06-29). EVERY build of EVERY component (source compile, artifact/package/installer/container-image production, codegen, asset render — for ANY language/platform: Go, Android/Gradle, desktop, web, native, firmware) MUST run INSIDE a specialized build container provisioned via the `digital.vasic.containers` submodule (§11.4.76) — NEVER on the bare host. The build containers MUST be DISTRIBUTED to the designated remote build host(s) (e.g. `thinker.local`) via the SAME containers-submodule distribution mechanism the infra uses (§11.4.76 Distributor / remote compose over SSH, §11.4.161 rootless), so the build EXECUTES on the remote build host (offloading the developer/main host); once the build completes the produced artifacts MUST be brought BACK to the originating main host (scp/rsync/volume copy) for use/flashing/distribution. Building outside a container, or on the bare host, is FORBIDDEN — a release blocker (the "works on my machine" / unreproducible-build class §11.4.76 exists to prevent). The build-host target + build-container definitions are config-injected (§11.4.28, never hardcoded in the submodule); a missing build-container capability is added by EXTENDING the containers submodule upstream (§11.4.74), never by an ad-hoc host build. Honest boundary (§11.4.6): the containerized+distributed build guarantees reproducibility + host-isolation + capacity offload — it does NOT replace the §11.4.40 full-suite retest, §11.4.108 four-layer artifact → runtime verification, or §11.4.38 installable-asset evidence (those still run against the brought-back artifact). Classification: universal (§11.4.17). Composes §11.4.76 (containers submodule — sole orchestration layer), §11.4.161 (rootless runtime), §11.4.74 (extend-don't-reimplement), §11.4.28 (config injection / decoupling), §11.4.24 (build-resource stats), §11.4.82 (iteration speedup — persistent caches in the container), §11.4.121 (no-commit-while-build-writes-artifacts), §11.4.38 (installable-asset evidence on the brought-back artifact), §11.4.108 (artifact → runtime verification), §12.6 (host memory ceiling — offloading the build preserves the main host). Propagation gate `CM-COVENANT-114-173-PROPAGATION` (literal `11.4.173`) + recommended gate `CM-CONTAINERIZED-DISTRIBUTED-BUILD` (every build runs via the containers submodule on the remote build host; a bare-host build is detected + FAILs) + paired §1.1 mutation. **Canonical authority:** constitution submodule

Constitution.md §11.4.173. Non-compliance is a release blocker. No escape hatch — no `--build-on-host` , `--skip-container-build` , `--local-build-ok` , `--no-distributed-build` , `--bare-host-build` flag.

§11.4.174 — Shared-host process-ownership verification mandate (User mandate, 2026-06-29). On any shared host where other projects/users/agents may run concurrently, every inspection of or action on a process, build, daemon, port, lock, or system-state signal MUST first positively VERIFY the target is OURS before diagnosing or acting. Attribute a process to us ONLY by a strong discriminator — cwd inside this project's tree, argv naming our scripts/artifacts/repo path, a PID we launched + recorded (§11.4.147), or a lock/port path under our tree — NEVER a loose `pgrep java/gradle/node` name match (it matches every project on the host). NEVER kill/signal/restart a process not positively ours (other projects' gradle/java/podman/zig are off-limits — operator-workload-safety, §12/§11.4.133); the §12.8 daemon-kill remediation MUST be scoped to OUR daemons. When our work waits on host contention from another project's process, surface it (§11.4.66/§11.4.101) — block-don't-break. Every `pgrep/ps` filter MUST exclude self-matches + other-project matches. Forensic anchor (FACT 2026-06-29): a 200%-CPU java+gradlew on the shared host was a separate `catalogizer` project's gradle test (cwd=Projects/catalogizer), nearly mis-attributed to our build. Composes §11.4.6/§11.4.66/§11.4.101/§11.4.147/§12/§12.8/§11.4.133. Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-174-PROPAGATION` (literal `11.4.174`) + recommended gate `CM-PROCESS-OWNERSHIP-VERIFIED` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.174. Non-compliance is a release blocker. No escape hatch — no `--assume-process-ours` , `--loose-pgrep-ok` , `--kill-any-matching-process` , `--skip-ownership-check` flag.

§11.4.176 — Conflict-free multi-track work-division + exactly-once claim registry + capability-aware deadlock-proof device-lock (User mandate, 2026-07-02). Two mandatory multi-track arbitration layers, conflict-free by construction: (A) exactly-once work-item/logical-group claim registry (extends §11.4.147/§11.4.116; disjoint file-scope §11.4.58 L3 + worktree/CoW §11.4.167; no codebase loss §9.2/§11.4.113/§11.4.84/§11.4.41); (B) capability-aware deadlock-proof device-lock (§11.4.119 single-owner; all-or-nothing + non-blocking + TTL-reap breaks all 4 Coffman conditions; append-only JSONL + atomic snapshot §11.4.116); (C) universal/decoupled (§11.4.17/§11.4.28/§11.4.35) + evidence-based resource-tuning (§11.4.24/§11.4.6, never guessed). Classification: universal

(§11.4.17). Gates CM-WORK-DIVISION-EXCLUSIVE-CLAIM + CM-DEVICE-LOCK-DEADLOCK-FREE + CM-COVENANT-114-176-PROPAGATION + paired §1.1 mutation. Composes §11.4.6 §11.4.17 §11.4.24 §11.4.28 §11.4.35 §11.4.41 §11.4.58 §11.4.84 §11.4.85 §11.4.103 §11.4.113 §11.4.116 §11.4.119 §11.4.147 §11.4.167 §12.6 §9.2 §11.4.176.

§12.11 — Maximal dynamic resource utilization for containerized builds (User mandate, 2026-07-03). The containerized build path (its OWN cgroup + own MemoryMax → OOM-kills ITSELF, never user.slice) MUST use the maximal SAFE fraction of host capacity, computed DYNAMICALLY per host: auto-detect nproc/MemTotal/RLIMIT_NPROC (never hardcode, §11.4.6), reserve explicit host headroom (cores+RAM), scale -j against BOTH the memory model AND the process rlimit (privileged nproc-raise for full -j; EAGAIN-safe -j otherwise), never break/bottleneck/wedge. NO fixed low -j for the containerized path. The §12.6 60% ceiling + bare-metal -j cap REMAIN, SCOPED to user.slice-resident work — §12.11 does NOT weaken §12.6, it scopes it. Classification: universal (§11.4.17). Gate CM-MAXRES-DYNAMIC-BUILD + CM-COVENANT-12-11-PROPAGATION + paired §1.1 mutation. Composes §12.1 §12.2 §12.3 §12.6 §12.10 §11.4.6 §11.4.24 §11.4.133 §9.2 §11.4.35 §12.11.

§11.4.177 — Developer-tooling project-decoupling + invocation-directory operation (User mandate, 2026-07-04). Compact summary: no project-specific script / hook / alias / binary may be wired (copied / symlinked / PATH-exported / aliased) into a global or shared developer-tooling PATH; shared tooling MUST be project-agnostic and operate on the INVOCATION directory (cwd / explicit path / config arg), NEVER a hardcoded project path; a project-specific helper lives in + runs from its own repo, a genuinely-reusable helper is promoted to the shared toolkit ONLY after being stripped of every project-specific assumption (paths, device serials, package names, region endpoints) + taking its target from the invocation context; a project script reachable on the global/shared PATH is a decoupling violation of equal severity to importing project-specific code into a shared submodule (§11.4.28); forensic FACT 2026-07-04 — a project cwd-hook symlinked into the global Claude-Toolkit PATH re-coupled every alias to one hardcoded checkout; honest boundary §11.4.6 — decoupling guarantees project-agnostic, not correct (still needs the tool's own tests). Classification: universal (§11.4.17). Composes §11.4.28/.29/.35/.11/.6. Propagation gate CM-COVENANT-114-177-PROPAGATION (literal 11.4.177) + recommended gate CM-TOOLING-PROJECT-DECOUPLED + paired §1.1 mutation. **Canonical authority:** constitution submodule

Constitution.md §11.4.177. Non-compliance is a release blocker. No escape hatch — no `--global-symlink-ok` , `--hardcode-project-path` , `--wire-into-shared-path` flag.

§11.4.178 — Track-qualified identity for parallel work streams (session-name collision ban) (User mandate, 2026-07-04). Compact summary: parallel work streams sharing hardware / a git backing store / a tmux-session namespace / a lock namespace / any single-owner resource MUST be addressed by a TRACK-QUALIFIED identity (`<project>__<track>__<role>`), NEVER a bare directory basename; multiple checkouts/clones/worktrees sharing basename `atmosphere` collapse into one session/lock/log key + cross-wire tmux targets, lock files, device claims, log dirs (observed collision 2026-07-04); every registry row / session name / lock path / log dir / device claim for a ≥ 2 -concurrent-claimant resource MUST carry the track-qualified key, a bare-basename identity for such a resource is a violation; honest boundary §11.4.6 — a unique identity prevents cross-wiring, it does not itself arbitrate ownership (that is §11.4.119 + **§11.4.176** (the multi-track work-division + exactly-once-claim + device-lock arbitration anchor)). Classification: universal (§11.4.17). Composes §11.4.111/.119/.29/.58 + **§11.4.176** (the multi-track work-division + exactly-once-claim + device-lock arbitration anchor). Propagation gate `CM-COVENANT-114-178-PROPAGATION` (literal `11.4.178`) + recommended gate `CM-TRACK-QUALIFIED-IDENTITY` + paired §1.1 mutation. **Multi-project-per-track layout clarification (2026-07-10):** a `/mnt/track<N>` track parent is NOT single-project — MULTIPLE projects coexist on the same track parents, each in its OWN project-named subdirectory `/mnt/track<N>/<project_name>/` (e.g. `/mnt/track1/atmosphere/` AND `/mnt/track1/helix_ota/` side by side); the per-track working copy of project P on track N is ALWAYS `/mnt/track<N>/<project_name>/` (canonical lowercase snake_case dir-name per §11.4.29), never the bare track root and never a shared basename; this is WHY the identity carries the `<project>` component (`<project>__<track>__<role>`) — so two coexisting projects' track-N sessions / locks / logs / device-claims / registry rows never collide; a project adopting the multi-track engine (§11.4.187) MUST create AND register its own `/mnt/track<N>/<project_name>/` per track and MUST NOT assume it owns the track root or that a track hosts only one project.

Canonical authority: constitution submodule Constitution.md §11.4.178. Non-compliance is a release blocker. No escape hatch — no `--bare-basename-ok` , `--skip-track-qualifier` , `--shared-session-name` flag.

§11.4.179 — Corruption-isolated parallel git streams (own-.git independent clones, not shared-common-dir worktrees) (User mandate, 2026-07-04).

Compact summary: when corruption-isolation is a requirement, parallel git work streams MUST each be an isolated repo with its OWN `.git` (own object store, own index, own lock namespace), NOT `git worktree` checkouts sharing one common `.git`; a shared common-dir is a single point of failure — one stale `index.lock` / `.commit_all.lock` / `.push_all.lock` freezes EVERY stream at once (observed 9-h all-track freeze 2026-07-04) + one corrupted object store loses every stream; each stream's `git rev-parse --git-common-dir` MUST resolve INSIDE its own tree, a common-dir resolving to a shared parent is a violation; where CoW/reflink disk-feasibility is the constraint use the §11.4.167 reflink-clone-with-own-`.git` path (btrfs/XFS/ZFS/APFS reflink shares extents on disk while keeping a per-stream `.git`, so isolation + disk-feasibility both hold); honest boundary §11.4.6 — own-`.git` contains lock/corruption blast radius, it does NOT remove per-repo stale-lock reaping (§11.4.180) nor ff-only no-force integration (§11.4.113). Classification: universal (§11.4.17). Composes §11.4.167/.119/.58/§9.2 + **§11.4.176** (the multi-track work-division + exactly-once-claim + device-lock arbitration anchor). Propagation gate `CM-COVENANT-114-179-PROPAGATION` (literal `11.4.179`) + recommended gate `CM-GIT-STREAMS-OWN-COMMON-DIR` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.179. Non-compliance is a release blocker. No escape hatch — no `--shared-common-git`, `--worktree-when-isolation-required`, `--skip-git-isolation` flag.

§11.4.180 — Commit/push (single-writer) wrappers MUST auto-reap provably-stale locks before acquiring (User mandate, 2026-07-04).

Compact summary: every commit / push / single-writer wrapper MUST, before acquiring its own lock, auto-reap any PROVABLY-stale git lock — one whose recorded holder PID is DEAD (`kill -0` fails), OR (no PID recorded) older than a defined threshold AND with no live wrapper/git process holding that repo; the reap covers the wrapper's own lock plus the git-internal existence-based locks (`index.lock`, `HEAD.lock`, `refs/**/*.*.lock`) git does NOT auto-release on holder death; it MUST NEVER remove a lock whose holder is ALIVE (removing a live-held lock corrupts a concurrent writer — §9.2); each reap decision (REAPED / KEPT + reason) is logged as captured evidence (§11.4.6 — liveness PROVEN via `kill -0`, never assumed); a wrapper that blocks indefinitely on a dead-holder lock (observed 9-h freeze 2026-07-04) is a violation; honest boundary §11.4.6 — reaping clears dead-

holder deadlock, it does NOT substitute for the own- .git isolation of §11.4.179. Classification: universal (§11.4.17). Composes §11.4.84/.88/.6/§9.2/§11.4.116/§11.4.179. Propagation gate `CM-COVENANT-114-180-PROPAGATION` (literal `11.4.180`) + recommended gate `CM-WRAPPER-STALE-LOCK-AUTO-REAP` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.180. Non-compliance is a release blocker. No escape hatch — no `--skip-stale-lock-reap`, `--force-remove-live-lock`, `--block-on-dead-holder` flag.

§11.4.181 — Consistent controlled feature-branch naming: one feature/logic-group ⇒ exactly ONE canonical branch name across the main repo + all owned submodules (User mandate, 2026-07-05). Compact summary: one feature / logical group of workable items MUST map to EXACTLY ONE canonical branch name used IDENTICALLY on the main repo AND every touched owned submodule — NEVER two differently-named branches for one feature/group, NEVER per-repo naming drift; (1) single canonical name per group minted ONCE per the §11.4.167 D scheme (`feature/<slug> branch + <prefix>-<base>-feat-<slug> tag, <slug> lowercase snake/kebab §11.4.29`), used identically across the main repo + every branched submodule (an untouched submodule stays on its base branch); (2) single source of truth — the canonical name is recorded ONCE in the §11.4.176 exactly-once claim registry / the §11.4.93 workable-items `logic_group → branch` mapping and LOOKED UP, never re-invented (creating a branch whose name ≠ the group's registered canonical name, or minting a 2nd name for a group that already has one, is a §11.4.6 no-guessing violation); (3) mechanical enforcement (§11.4.75) — recommended gate `CM-BRANCH-NAME-CONSISTENCY` FAILs on an unregistered feature-branch name / two divergent names for one group / a submodule branch name diverging from the group's main-repo canonical name, paired §1.1 mutation (register one group under two divergent names → gate FAILs); (4) risk-free reconciliation (§9.2/§11.4.113/§11.4.84) — a mis-named/duplicate branch is reconciled by MERGING its commits onto the canonical branch (union, no commit lost, no `-s ours /reset`), ff-only push to all upstreams (NEVER force-push), retiring the mis-named branch only after the merge is confirmed on all mirrors. Classification: universal (§11.4.17). Composes §11.4.167/.176/.178/.179/.28/.29/.93/.113/.84/.6/.75/§9.2. Propagation gate `CM-COVENANT-114-181-PROPAGATION` (literal `11.4.181`) + recommended gate `CM-BRANCH-NAME-CONSISTENCY` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.181. Non-compliance is a

release blocker. No escape hatch — no `--allow-divergent-branch-names` , `--skip-branch-registry` , `--rename-branch-freely` , `--per-repo-branch-name-OK` , `--two-names-one-feature-OK` flag.

§11.4.182 — Track+branch work-stream identity label (T<N>/<branch>) on every agent/subagent label + every operator-facing work-stream reference (User mandate, 2026-07-05). Compact summary: EVERY agent / subagent / work-stream label AND every operator-facing reference to a work stream MUST START with a (T<N>/<branch> - <alias>) prefix identifying the track + git branch + the claude-toolkit alias in use (e.g. (T1/main - claude1) ATM-312 MPV drm_prime investigation , (T2/feature/mistiq-vader - deepseek) SPK-XXX ...) so parallel-track work (/mnt/track1..4 , each on its own branch, each driven by a distinct toolkit alias) is never ambiguous; (1) universal label form on every agent-dispatch description, status report, resume/handoff doc, progress line; (2) DETERMINISTIC derivation never guessed (§11.4.6) — <N> from the checkout path /mnt/track<N>/... (honest ? off-track, never a guessed number), <branch> from git rev-parse --abbrev-ref HEAD , <alias> from CLAUDE_CONFIG_DIR basename .claude-<alias> → <alias> (honest ? when unset/non-matching, never a guessed name) (reference labeler scripts/multitrack/track_branch_label.sh); (3) mechanical enforcement (§11.4.75 / §11.4.109-class) — a PreToolUse guard hook (constitution/scripts/hooks/guard-track-branch-label.sh , inherited BY REFERENCE per §11.4.177, never copied) BLOCKS any Agent/Task/TaskCreate dispatch whose description (or subagent label) does not START with ^\([0-9]+\^[^)]+\) , non-agent tools untouched; pre-build gate CM-TRACK-BRANCH-LABEL (labeler+hook exist/exec/parseable, hook wired in settings, convention doc exists) + propagation gate CM-COVENANT-114-182-PROPAGATION (literal 11.4.182); (4) honest boundary (§11.4.6) — a ? in the track OR the alias field is a valid honest label, never a bluff, and the enforced numeric-track format surfaces (blocks) an off-track dispatch rather than mislabeling it. Classification: universal (§11.4.17). Composes §11.4.178 (track-qualified identity — the label-prefix instantiation applied to every agent/subagent label + operator-facing reference) / §11.4.29 / §11.4.75 / §11.4.109 / §11.4.177 / §11.4.6 / §11.4.58 / §11.4.103 / §11.4.119 / §11.4.116 / §11.4.147. Propagation gate CM-COVENANT-114-182-PROPAGATION (literal 11.4.182) + recommended gate CM-TRACK-BRANCH-LABEL + paired §1.1 mutation. **Canonical authority:** constitution submodule Constitution.md §11.4.182. Non-compliance is a

release blocker. No escape hatch — no `--skip-track-label` , `--unlabeled-agent-OK` , `--bare-work-stream-reference` , `--no-track-branch-prefix` , `--label-from-memory` flag.

§11.4.183 — Maximal multi-agent utilization per work-stream + full-constitution-application + zero-bluff mandate (User mandate, 2026-07-07).

Compact summary: every track / work-stream MUST maximize multi-agent working approaches — subagent-driven development (§11.4.20/§11.4.70), independent review agents (§11.4.125/§11.4.142/§11.4.165/§11.4.134), parallel background streams with auto-backfill (§11.4.58/§11.4.103), and every other applicable agent form — so every track produces the MOST completed work at the HIGHEST quality in the LEAST time; every track MUST apply the ENTIRE constitution (every rule, mandatory constraint, guide, and derived technology — NOTHING ignored, skipped, avoided, or deemed less-relevant); there MUST NEVER be false / faulty / untested / unverified results or bluff of any kind anywhere (§11.4 anti-bluff covenant, captured evidence §11.4.5/§11.4.69/§11.4.107); honest boundary §11.4.6 — "maximize agents" is bounded by §12.6 60% memory + §11.4.58 parallel-agent cap + genuine parallelizability + §11.4.119 single-resource-owner, spawning contending/redundant agents is NOT maximization, the bar is maximum USEFUL parallel throughput not agent count. Classification: universal (§11.4.17). Composes §11.4.20/.58/.70/.92/.94/.97/.103/.119/.125/.126/.134/.142/.165/§12.6. Propagation gate `CM-COVENANT-114-183-PROPAGATION` (literal `11.4.183`) + recommended gate `CM-MAX-AGENTS-PER-TRACK` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.183. Non-compliance is a release blocker. No escape hatch — no `--single-agent-OK` , `--skip-review-agents` , `--partial-constitution-application` , `--serialise-parallelisable-work` flag.

§11.4.184 — Mandatory SonarQube static-analysis CLI + local tooling installed and PATH-discoverable (User mandate, 2026-07-06).

Compact summary: every project/host that runs static analysis MUST have the SonarQube scanner CLI (`sonar-scanner`) installed AND durably discoverable on the system PATH via the shell rc (`.bashrc` / `.zshrc`) — exactly like the Firebase/GitHub/GitLab CLIs — PLUS the shared `constitution/scripts/sonarqube/` tooling (`sonarqube_install_check.sh` / `sonarqube_lib.sh` / `sonarqube_run_scan.sh` / `sonarqube_container.sh` + `compose/`) consumed BY REFERENCE (§11.4.28 decoupled + §11.4.177 project-agnostic + §11.4.80-style inherited, NEVER copied per project); (1) installed + PATH-durable — the

scanner resolves in every fresh shell via a durable rc export, not an ephemeral session-only PATH; (2) shared tooling present + GREEN — `sonarqube_install_check.sh` exits 0 (scanner + rootless podman §11.4.161 + podman-compose + curl + ES-ready `vm.max_map_count`) before any scan-dependent work; (3) anti-bluff §11.4.6 — "installed/configured" PROVEN by `command -v sonar-scanner` resolving AND install-check exit 0, NEVER assumed, a claimed-installed-but-absent CLI is a §11.4 tooling-layer bluff; (4) rootless §11.4.161 — the local SonarQube server runs via rootless podman, never rootful docker/sudo; honest boundary §11.4.6 — names the SonarQube SCANNER CLI + shared local-server tooling, and per §11.4.166 (former universal Semgrep static-analysis mandate REPEALED) SonarQube is a DISTINCT operator-mandated tool not a Semgrep re-instatement, mandating TOOLING availability not a scan on every build. Classification: universal (§11.4.17). Composes §11.4.28/.36/.74/.75/.109/.161/.166/.177/.6. Propagation gate `CM-COVENANT-114-184-PROPAGATION` (literal `11.4.184`) + recommended gate `CM-SONARQUBE-CLI-INSTALLED` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.184. Non-compliance is a release blocker. No escape hatch — no `--skip-sonarqube-install` , `--sonar-cli-optional` , `--ad-hoc-sonar-path` , `--rootful-sonar-OK` flag.

§12.12 — Process/thread-limit (RLIMIT_NPROC) awareness for parallel subagent/multi-process work (User mandate, 2026-07-07). Compact summary: heavy parallel subagent/multi-process work is bounded by the OS per-user process/thread limit `ulimit -u` (`RLIMIT_NPROC` , Linux counts ALL threads+processes of the real UID, not per-process); the ambient tooling fleet (MCP servers, tokio/V8/libuv worker runtimes, `mongod` / `chrome` / `prisma` -class MCP servers, local LLM servers) consumes a large FIXED fraction before project work starts — FORENSIC FACT (2026-07-07): a host at `ulimit -u 4096` sat at ~3900-4005 threads from the ambient fleet alone (<150 headroom); a 4th Go-heavy subagent (`GOMAXPROCS=nproc` spawns dozens of threads) hit `EAGAIN` / `failed to create new OS thread (errno=11)` and crashed tooling; stopping 3 non-essential MCP servers freed ~1345 threads (3827 → 2482). Mandate: BEFORE scaling parallel subagents/processes, check thread headroom (`ulimit -u soft + ulimit -Hu hard + live ps -L --no-headers -u "$USER" | wc -l`) and treat exhaustion as a §12 host-safety event that YIELDS UNCONDITIONALLY — the ORTHOGONAL second exhaustion axis alongside §12.6's memory ceiling (abundant RAM ≠ abundant thread budget); signals `EAGAIN` / `fork: retry` /

failed to create new OS thread / cannot allocate memory on fork/clone (distinct from OOM, §11.4.6 no-guessing — never conflate); low headroom ⇒ SERIALIZE, avoid `GOMAXPROCS=nproc` -class spikes, never dispatch blind (a subagent hitting `EAGAIN` mid- git push is a §9.2 data-safety risk). Three raising techniques: (a) no-restart `prlimit --pid <PID> --nproc=<soft>:<hard>` (root, e.g. `su -c`) — new children of a running process inherit the raise; (b) persistent `/etc/security/limits.conf` (`<user> hard/soft nproc <N>`) or `systemd` `LimitNPROC=<N>`, next login; (c) self-raise `ulimit -u <N>` up to the existing hard ceiling within the current session — `su -c "ulimit -u N"` sets the limit ONLY inside that transient subshell, NOT the already-running session. Freeing threads (stopping non-essential MCP servers) REQUIRES operator authorization (§11.4.122, no autonomous tooling removal); CAUTION `pkill -f` / `pgrep -f` match the killer's OWN command line — always exclude the caller's `$$` / `$PPID` (and conductor PID) from any cmdline-pattern kill list to avoid self-kill.

Classification: universal (§11.4.17). Composes §12/.6/.7/§11.4.58/.101/.103/.122/ §9.2. Propagation gate `CM-COVENANT-12-12-PROPAGATION` (literal `12.12`) + recommended gate `CM-NPROC-HEADROOM-CHECK` + paired §1.1 mutation.

Canonical authority: constitution submodule `Constitution.md` §12.12. Non-compliance is a release blocker. No escape hatch — no `--ignore-thread-limit`, `--dispatch-blind`, `--skip-headroom-check`, `--autonomous-tooling-kill`, `--assume-ulimit-scope` flag.

§11.4.185 — Manual QA-team testing as the mandatory FINAL confirmation of done (User mandate, 2026-07-07). Compact summary: no set of features / scope of work / release may be considered fully completed until it has received MANUAL TESTING confirmation by the project's QA team as the FINAL confirmation step — everywhere and always, never forgotten; every automated gate (§11.4 anti-bluff covenant, §11.4.5/§11.4.69 captured evidence, §11.4.52 autonomous validation, §11.4.40 full-suite retest, §11.4.108 four-layer verification, §11.4.132 risk-ordered validation) is NECESSARY but NOT SUFFICIENT — manual QA is the FINAL human sufficiency gate; ADDS to (never weakens) §11.4.52's mandatory autonomous-validation-path requirement — a feature needs BOTH the autonomous path AND the manual-QA final confirmation, never one substituting the other; pipeline: autonomous-GREEN → build → flash/deploy → autonomous on-device validation → hand off to QA team → ONLY after QA-team manual confirmation is the scope fully-completed / tag-eligible; the §11.4.126 autonomous-loop release-tag terminal condition now REQUIRES this confirmation; honest boundary §11.4.6 —

the agent MUST hand off + WAIT, never self-certify/simulate/infer the manual step, while continuing every other non-blocked item in parallel per §11.4.87/.94/.97/.101 (the wait parks one scope, not the loop). Classification: universal (§11.4.17). Composes §11.4/.5/.6/.40/.52/.69/.108/.126/.129/.130/.132. Propagation gate `CM-COVENANT-114-185-PROPAGATION` (literal `11.4.185`) + recommended gate `CM-MANUAL-QA-FINAL-CONFIRMATION` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.185. Non-compliance is a release blocker. No escape hatch — no `--skip-manual-qa`, `--automation-suffices`, `--self-certify-manual-step`, `--assume-qa-confirmed`, `--autonomous-only-completion` flag.

§11.4.186 — Anti-divergence enforcement: cross-document consistency is a mandatory-before-export/commit gate, never an after-the-fact audit (research-derived, 2026-07-08). Compact summary: any project maintaining more than one representation of the same tracked data (a delivery plan + its MVP brief + a workable-items DB + their summaries + `.md` / `.html` / `.pdf` / `.docx` exports) MUST enforce cross-document CONSISTENCY as a deterministic gate that runs BEFORE any export render, any doc/DB sync `verify`, and any doc-set commit — NEVER as an after-the-fact human audit that runs only when an operator happens to ask (the forensic FACT: the same NanoKVM/phone-as-screen feature appeared in three plan clusters — one shared ticket SPK-596 across SIX rows split over TWO releases, and a carve-out due to ship 2026-08-31 whose own prerequisite does not START until 2026-10-05 — and NO gate caught it; only an operator's manual read of `Plan_v9.0.xlsx` surfaced it, after the divergent bytes had already exported). (1) **One no-divergence verdict, three seams** — a single deterministic PASS/FAIL/SKIP verdict (built by composing a cross-document integrity validator with the workable-items-DB internal `validate`) gates all three write-seams (export / sync-verify / commit); a FAIL at any seam refuses that seam, naming the exact offending records (§11.4.6 actionable). (2) **Five decidable check families** — DEDUP (no two records describe one feature with divergent timeline/status/release, keyed on ticket OR normalised `(subject, scope)`, NOT bare subject substring), TIMELINE (start ≤ deadline; no deadline-before-dependency; no dependency cycle; GATED exempt-but-marked), CROSS-DOC (the same item across every representation carries identical timeline/status/type against a NAMED authoritative source), INTEGRITY (no orphan refs; Status ↔ Type §11.4.33; location ↔ status), STRUCTURAL (required columns; ID uniqueness+monotonicity §11.4.54). (3) **Drift-proof trigger** — a §11.4.86 sha256-of-sorted-members fingerprint of the

authoritative inputs re-arms the gate on ANY input change, so a stale verify cannot pass on changed data. (4) **Self-validated analyzer (§11.4.107(10))** — the validator ships golden-good (MUST PASS) + one golden-bad per family (MUST FAIL, pinpointing the offender) + a negative-control (distinct same-subject tasks that MUST PASS — the false-positive guard); a validator that passes its golden-bad, or fails golden-good/the negative-control, is itself a §11.4 bluff and a release blocker. (5) **Reusable + decoupled (§11.4.28/31)** — packaged as a depth-1 reusable engine under the constitution submodule per the §11.4.28(C) carve-out with a `helix-deps.yaml` ; project-specific doc-sets + authoritative-source bindings live in a consumer-owned checkset (data, never engine code). (6) **Supersedes ad-hoc audits** — the per-cycle `*_INTEGRITY_FINDINGS.md` / `RECONCILIATION_*.md` after-the-fact audit docs are retired in favour of the gate; honest boundary §11.4.6 — the gate proves internal CONSISTENCY (no record diverges, no timeline is incoherent, every ref resolves), NOT plan CORRECTNESS (achievability/date-rightness stay operator/management decisions the gate SURFACES, never MAKES). Classification: universal (§11.4.17). Composes §11.4.6/.12/.33/.44/.50/.53/.54/.60/.65/.73/.75/.85/.86/.91/.93/.95/.106/.107/.108/.110/.148/.176. Propagation gate `CM-COVENANT-114-186-PROPAGATION` (literal `11.4.186`) + recommended gate `CM-DOC-INTEGRITY-VALIDATION` (5 invariants: validator present+executable; consumer checkset present; the no-divergence gate wired into each of the three seams; fingerprint sidecar fresh; `selfcheck` golden-good/golden-bad/negative-control wired into meta-test) + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.186. Non-compliance is a release blocker. No escape hatch — no `--skip-divergence-gate` , `--audit-after-export-OK` , `--cross-doc-check-not-applicable` , `--unvalidated-analyzer-OK` , `--divergent-commit-OK` flag.

§11.4.187 — Automatic multi-track ruler orchestration (User mandate, 2026-07-04; landed 2026-07-09). Compact summary: every project incorporating this constitution MUST ship its multi-track parallel-development orchestration as a UNIVERSAL, AUTOMATIC, OUT-OF-THE-BOX, INHERITED capability — not a per-project reimplementation — via a single conductor session that programmatically SPAWNS, DRIVES, RESUMES, MONITORS, and CONTROLS per-track headless workers. (1) **Headless spawn primitive** — workers launch via native headless `claude -p --output-format stream-json --verbose` ; the worker's `session_id` is captured from the FIRST `init` event (never invented);

success/failure is read from the terminal event's `result.is_error` field, NEVER the process exit code (a CLI can exit non-zero on a successful `result`, or zero on a failed one — exit-code-as-success-signal is a §11.4.6 guess); continuation of an existing worker uses `--resume <session_id>` invoked from the SAME cwd + env the spawn used (the storage/resume-collision gotcha). (2) **Per-subscription auth + rate-limit fallback** — every spawn `unset` s any ambient `ANTHROPIC_API_KEY` (the precedence trap that silently routes a "subscription" worker onto pay-per-token API billing) and sets the selected alias's `CLAUDE_CODE_OAUTH_TOKEN` / `CLAUDE_CONFIG_DIR`; on a captured rate-limit signature (`isApiErrorMessage + apiErrorStatus:429`, quota-exhausted vs transient distinguished by evidence, never guessed) the ruler REBINDS the affected track onto the next healthy alias (cross-track reuse permitted), falls back to a provider alias where configured, and bounded-parks (§11.4.101) rather than hanging when every alias is cooled. (3) **Crash-resilient ruler self-supervisor** — the ruler's OWN state (per-track alias/session_id/worktree/last-verdict) is persisted durably (write-temp-then-rename, NEVER tmpfs) so an external watchdog detects a dead ruler PID and re-hydrates every track's binding from the durable state + the §11.4.116 event stream — no work is silently lost to a ruler crash (§11.4.147 at the ruler layer). (4) **Idempotent bootstrap + auto-install + conductor-stays-home** — one idempotent bootstrap command wires the whole system out of the box (cwd-hook install, config validation, reconciliation), auto-invoked on every constitution pull via the §11.4.164 `post_update_hook.sh` skill-register seam; the `conductor`: -designated alias/session is NEVER bound into a track worktree (the resolver returns nothing for it) so the conductor always runs from the main checkout. (5) **Project-agnostic engine, consumer-owned data** — the entire mechanism (config loader, alias → worktree resolver, cwd-hook, bind/fallback/cooldown orchestrator, session launcher, rate-limit monitor, supervisor) lives at `constitution/scripts/multitrack/`, inherited BY REFERENCE (§11.4.28(B)/§11.4.177), carrying NO project literal; every consuming project supplies its own per-host `config/multitrack/<hostname>.yaml` (tracks/mounts/serials/aliases/conductor) as DATA, never as an engine edit. Anti-bluff (§11.4/§11.4.108): every piece carries four-layer coverage (pre-build gate + on-host test + paired §1.1 mutation + captured evidence per §11.4.69), and each fix declares a runtime signature verified on a clean host — never a claim the mechanism "should work." Honest boundary (§11.4.6): this anchor mandates the MECHANISM; genuine per-subscription quota-isolation proof requires live tokens exercising real rate-limit

boundaries (an operator-scheduled live acceptance window) — a mock-driven GREEN proves the spawn/resume/fallback CONTRACT, never live-quota behaviour, and the two are never conflated or substituted for each other. Classification: universal (§11.4.17). Composes §11.4.20 (subagent-driven-by-default) / §11.4.28 (submodule decoupling + dependency layout) / §11.4.58 (parallel-development PWU pipeline) / §11.4.70 (subagent-driven execution default) / §11.4.101 (autonomous-decision-over-blocking — the bounded-park-on-all-cooled behaviour) / §11.4.103 (continuous parallel-stream routine) / §11.4.111 (resolve-by-stable-name — config by serial, never by index) / §11.4.116 (real-time conductor ↔ framework sync channel) / §11.4.119 (single-resource-owner partitioning) / §11.4.126 (default autonomous-loop mode) / §11.4.128 (always-on device-recording — the analogous host-safety discipline) / §11.4.131 (standing session-resumption file) / §11.4.147 (crashed-agent respawn-until-complete — the ruler-self-supervisor is its ruler-layer instantiation) / §11.4.164 (constitution auto-propagation hook seam) / §11.4.167 (feature work-stream lifecycle) / §11.4.176 (multi-track work-division + exactly-once claim + device-lock) / §12.6/§12.8 (host-safety memory + concurrency ceilings — the spawn/launch pool is budget-gated). Propagation gate `CM-COVENANT-114-187-PROPAGATION` (literal `11.4.187`) + recommended gates `CM-MULTITRACK-ENGINE-IN-CONSTITUTION` (the generic ruler scripts present under `constitution/scripts/multitrack/`, carrying no project literal, `bash -n + sh -n` clean per §11.4.67) / `CM-MULTITRACK-AUTOINSTALL-WIRED` / `CM-MULTITRACK-BOOTSTRAP-IDEMPOTENT` / `CM-MULTITRACK-BIND-ON-START` / `CM-MULTITRACK-FALLBACK-MONITOR` / `CM-MULTITRACK-HEADLESS-SPAWN` / `CM-MULTITRACK-RULER-SUPERVISOR` + paired §1.1 mutation (strip the literal → propagation gate FAILs; re-introduce a project literal into the engine, or downgrade a rate-limit rebind to exit-code-based success detection → the corresponding mechanism gate FAILs). **Canonical authority:** constitution submodule `Constitution.md` §11.4.187. Non-compliance is a release blocker. No escape hatch — no `--skip-multitrack`, `--manual-multitrack-only`, `--no-auto-install`, `--per-project-reimpl-OK`, `--exit-code-is-success-signal` flag.

§11.4.188 — Regular main → feature merge cadence: every track / feature-branch MUST FREQUENTLY merge canonical `main` into itself DURING the work, never only at the end (User mandate, 2026-07-09). Compact summary: every long-lived feature branch AND every parallel-development track (§11.4.58/ §11.4.103 parallel streams; §11.4.167 feature work-streams; §11.4.176/§11.4.178/

§11.4.179 multi-track) MUST regularly `git merge origin/main` (the canonical trunk/master) INTO its own branch THROUGHOUT the work — NOT only at the end — so no branch ever drifts far from main and the eventual back-merge stays a small, low-conflict, low-risk operation. GENERALISES §11.4.167(D) (trunk-merged-INTO-every-BIG-feature-STREAM, at minimum after every trunk tag / daily, `git merge` never rebase, stale-if-not-synced) from the big-work-item feature-stream lifecycle to EVERY feature branch on EVERY track — a small/medium feature branch, or any `git worktree` /CoW track that never earned a full §11.4.167 stream, is STILL bound to this cadence (the §11.4.167(D) discipline promoted to the general standalone rule; §11.4.167(D) remains its big-item specialisation). (1) **WHAT** — fetch the latest canonical `main` / `master` and merge it INTO the working branch, keeping the branch continuously close to trunk; a branch that has NOT integrated trunk within the cadence is flagged STALE (the exact long-lived-branch conflict-storm + integration-risk anti-pattern this forbids). (2) **CADENCE** — merge trunk in FREQUENTLY: after EVERY trunk tag, at minimum DAILY while the branch is live, AND before starting any significant new chunk of work on the branch — MANY SMALL merges, never one big deferred end-of-branch merge (small frequent merges minimise the divergence + conflict surface). (3) **HOW (safe-by-construction)** — MERGE, NEVER rebase a shared/tagged/pushed branch (no commit-loss, no history rewrite — §9/§11.4.113); FETCH-FIRST every remote before the merge (§11.4.37/§11.4.71 — remote state is unknowable without a fetch, §11.4.6); take a §9.2 hardlinked pre-op backup before any LARGE / conflict-heavy / risky merge; resolve every conflict CAREFULLY — ZERO conflict markers committed, ZERO file silently dropped, union of both sides preserved (§11.4.41 step 3 / §9 no-loss); integrate ONLY when the branch is QUIESCENT (§11.4.84 — no in-flight mutation gate, no unaccounted uncommitted work); prefer running the merge + its verification in the BACKGROUND parallel to the main stream so it never stalls other work (§11.4.88/§11.4.89/§11.4.103); NEVER force-push the merged result (§11.4.113 — the merged branch fast-forwards its own mirrors). (4) **CONSISTENCY + ANTI-BLUFF** — after EVERY merge the branch MUST be PROVEN consistent, never ASSUMED: a pre-build smoke gate runs post-merge (§11.4.42 step-3 smoke) and is GREEN, a `git grep` for conflict markers (`<<<<<<<` / `=====` / `>>>>>>>`) returns EMPTY, and no commit present pre-merge on either side is lost — captured evidence per §11.4.5/§11.4.69; a merge claimed "clean" without the post-merge smoke + marker-scan + no-lost-commit check is a §11.4/§11.4.1 bluff at the integration layer. (5) **HONEST BOUNDARY (§11.4.6)** — regular trunk-integration keeps the branch MERGEABLE + low-conflict;

it does NOT by itself prove the branch's own work is correct (each change still crosses §11.4.108 runtime-signature + §11.4.40 retest), and it is NOT the approval-gated back-merge TO trunk (that stays §11.4.167(I)/§11.4.40/§11.4.41 no-merge-until-approved). Classification: universal (§11.4.17) — the consuming project supplies its canonical-trunk ref name, cadence unit, post-merge smoke-gate command, and per-track worktree/CoW mechanism per §11.4.35. Composes §9/§9.2/§11.4.6/§11.4.17/§11.4.37/§11.4.41/§11.4.42/§11.4.71/§11.4.84/§11.4.88/§11.4.89/§11.4.103/§11.4.113/§11.4.167/§11.4.176/§11.4.178/§11.4.179/§11.4.181. Propagation gate `CM-COVENANT-114-188-PROPAGATION` (literal `11.4.188`) + recommended gate `CM-REGULAR-MAIN-MERGE-INTO-FEATURE` (every live feature branch / track has integrated the canonical trunk within the declared cadence; its last integration is a real `git merge` not a rebase; the post-merge smoke + zero-conflict-marker + no-lost-commit checks passed) + paired §1.1 mutation (let a live branch go stale past the cadence, OR record a rebase of a shared branch in place of a merge, OR commit a conflict marker → the gate FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.188. Non-compliance is a release blocker. No escape hatch — no `--skip-main-merge`, `--merge-only-at-end`, `--rebase-shared-branch`, `--defer-trunk-sync`, `--no-post-merge-smoke`, `--force-push-after-merge`, `--big-items-only-trunk-sync` flag.

§11.4.189 — Most-reopened cases get extra-depth live-testing scrutiny FIRST (User mandate, 2026-07-10). Verbatim operator mandate: "All live testing MUST pay extra attention to cases which have been reopened numerous times and those cases in depth retested, investigated, fully validated and verified with real physical results and no bluff of any kind anywhere!" All LIVE TESTING (on-device / on-target runtime validation, deploy-and-observe) MUST give EXTRA-DEPTH retest + in-depth investigation + FULL validation/verification with REAL PHYSICAL captured evidence (§11.4.5/§11.4.69/§11.4.107 — captured audio/video/sysfs/dumpsys/sink-side/runtime-signature) and NO bluff of any kind ANYWHERE to the cases that have been REOPENED THE MOST TIMES (highest §11.4.55 reopens-count) — the empirically-most-fragile set gets the DEEPEST live scrutiny, and it gets it FIRST. (1) **WHAT** — for the highest-reopens-count cases, live testing is NOT a single pass: it re-runs the case at EXTRA DEPTH (more iterations per §11.4.50 deterministic consistency, wider edge-case + regression coverage), RE-INVESTIGATES the underlying defect per §11.4.102 systematic-debugging + §11.4.146 reproduce-first, and CONFIRMS with rock-solid captured physical evidence per §11.4.123 — never

metadata-only / config-only / absence-of-error / grep-without-runtime (§11.4/§11.4.1), never a false result. (2) **ORDER** — the most-reopened set is scrutinised FIRST, ahead of the rest of the live-test suite (the §11.4.132 risk-descending discipline, with most-reopened the load-bearing key). (3) **KEY** — priority is keyed off the §11.4.55 reopens-count (the consuming project supplies its reopens-count source per §11.4.35, e.g. the §11.4.93 workable-items DB `reopens_count`); a high reopens-count is the strongest empirical fragility signal (§11.4.132(d)). (4) **PROOF** — each most-reopened case's live PASS cites its captured-evidence artefact path (§11.4.69/§11.4.107) on a CLEAN/fresh deployment (§11.4.108 runtime-signature), with the §11.4.115 RED → GREEN polarity flip captured where a permanent regression guard exists (§11.4.135). STRENGTHENS/REFINES §11.4.132 (risk-ordered validation factor (d) most-reopened — §11.4.189 promotes most-reopened into its OWN EXTRA-DEPTH mandate at the LIVE-TEST layer, not merely one ranking factor among four) + §11.4.55 (reopens-count as the priority key) + §11.4.129/§11.4.130 (validate-the-just-fixed / highest-risk-first after redeploy). Classification: universal (§11.4.17) — references no project-specific hardware; the consuming project supplies its reopens-count source per §11.4.35. Composes §11.4.5/§11.4.6/§11.4.55/§11.4.69/§11.4.107/§11.4.108/§11.4.115/§11.4.129/§11.4.130/§11.4.132/§11.4.135/§11.4.146. Propagation gate `CM-COVENANT-114-189-PROPAGATION` (literal `11.4.189`) + recommended gate `CM-MOST-REOPENED-EXTRA-DEPTH-LIVE-TEST` (the live-test suite orders + extra-depth-scrutinises the highest-reopens-count cases FIRST, each with captured physical evidence on a clean deployment) + paired §1.1 mutation (drop the most-reopened extra-depth pass, OR run it AFTER the rest of the suite, OR PASS a most-reopened case on metadata-only evidence → the gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.189. Non-compliance is a release blocker. No escape hatch — no `--skip-most-reopened-depth` , `--reopened-same-as-any` , `--live-test-any-order` , `--metadata-pass-for-reopened-OK` , `--defer-reopened-scrutiny` flag.

§11.4.190 — Website engineering-quality mandate: every project website MUST be fully responsive + completely SEO-optimized + uniquely OpenDesign-authored + bleeding-edge enterprise-quality, each PROVEN with captured evidence (User mandate, 2026-07-10). Verbatim operator mandate: "every website we create inside our projects MUST BE fully responsive - working on all browsers, all operating systems, all platforms and all devices and screen

sizes! Every Website MUST BE completely SEO optimized! Layouts and templates done MUST BE unique and everything fully designed from OpenDesign! Make sure this applies to our project Website as well! Everything must be bleeding-edge enterprise quality, stunning visual look and feel, innovative, creative and unique! There MUST BE no false or faulty result or bluff of any kind anywhere!" Every website / web-UI surface a project ships — INCLUDING this project's OWN website — MUST satisfy AND PROVE with captured physical evidence (§11.4.5/§11.4.69/ §11.4.107/§11.4.170 — never a bare claim, §11.4.6) ALL of: **(A) FULL RESPONSIVENESS** — correct, usable rendering across ALL major browser engines (Chromium/Firefox/WebKit), all operating systems, all device classes (desktop/laptop/tablet/phone/large-display) and the full range of screen sizes/ orientations/pixel-densities; NO broken layout / overflow / clipping / overlap / off-screen or unusable control at ANY breakpoint. **(B) COMPLETE SEO OPTIMIZATION** — semantic HTML + correct document outline; per-page unique `<title>` + meta description; Open Graph + Twitter cards; canonical URLs; structured data (schema.org / JSON-LD); `robots` + XML sitemap; crawlable / indexable markup; WCAG AA accessibility + Core Web Vitals performance (both feed ranking). **(C) UNIQUE OpenDesign-AUTHORED templates/layouts** — EVERY layout + template designed FROM the OpenDesign design-token system (§11.4.162), NOT generic / off-the-shelf / templated; visually unique, innovative, creative, brand-distinct. **(D) BLEEDING-EDGE ENTERPRISE VISUAL QUALITY** — stunning, polished, professional look-and-feel in BOTH light AND dark themes (§11.4.162 tokens). **(E) ANTI-BLUFF PROOF** — NONE of (A)–(D) may be CLAIMED without captured physical evidence: responsiveness PROVEN by device-independent host-rendered screenshots across the real breakpoint × engine matrix + an OCR/vision layout oracle (no overlap / label-over-label / clip / overflow / off-screen); SEO PROVEN by an automated audit (Lighthouse SEO + structured-data validation + meta / OG / sitemap / robots presence) meeting a defined score floor; visual quality PROVEN by the §11.4.170 device-independent host-rendered pixel proof per screen × state × {light,dark}; uniqueness PROVEN by OpenDesign-token provenance. A website reported responsive / SEO-optimized / unique / enterprise-quality WITHOUT these captured proofs is a §11.4 PASS-bluff at the web-UI layer. §11.4.190 is the strict WEB-quality expansion of §11.4.162 (OpenDesign — it applies §11.4.162's design-token mandate to the full responsive + SEO + enterprise-quality web surface) and binds §11.4.170's device-independent host-rendered pixel proof as the correctness oracle; sibling of §11.4.168 (exported-document independent visual validation) at the web-UI layer. Classification:

universal (§11.4.17) — references no project-specific hardware; the consuming project supplies its concrete breakpoint × engine matrix, SEO score-floor, and OpenDesign token set per §11.4.35. Composes §11.4.5/§11.4.6/§11.4.69/§11.4.107/§11.4.162/§11.4.168/§11.4.170. Propagation gate `CM-COVENANT-114-190-PROPAGATION` (literal `11.4.190`) + recommended gates `CM-WEBSITE-RESPONSIVE-PROVEN` (host-rendered screenshots across the breakpoint × engine matrix + layout oracle — no overlap/clip/overflow/off-screen) / `CM-WEBSITE-SEO-OPTIMIZED` (automated SEO audit meets the score floor + structured-data / meta / OG / sitemap / robots present) / `CM-WEBSITE-OPENDSIGN-UNIQUE` (every layout/template's OpenDesign-token provenance, no generic/off-the-shelf template) + paired §1.1 mutation (report a website responsive/SEO/unique/enterprise-quality with NO captured proof, OR ship a generic non-OpenDesign template, OR skip the light/dark host-rendered pixel proof → the recommended gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.190. Non-compliance is a release blocker. No escape hatch — no `--skip-responsive-proof`, `--seo-optional`, `--generic-template-OK`, `--visual-quality-later`, `--claim-without-render` flag.

§11.4.191 — Work-to-track/branch binding enforcement: a logic-group's work can NEVER be committed OR dispatched onto the wrong track/branch (research-derived, 2026-07-10). Compact summary: every logical group of workable items binds to EXACTLY ONE canonical (`track`, `branch`) recorded ONCE in the single source of truth (the §11.4.93 workable-items DB `logic_groups.destination` = authoritative branch + a `canonical_track` column + a `group_paths` file-scope manifest), and NO change belonging to a group may LAND on, nor be DISPATCHED to, a track/branch other than that group's canonical one — the file-scope generalisation of §11.4.181/§11.4.182 from branch-NAME consistency to work-PLACEMENT consistency (forensic FACT 2026-07-10: `srcpy/Remote-app` work — group `mistiq-vader`, canonical `feature/mistiq-vader` /T2 — was committed onto `main` /T1; §11.4.181's `guard-branch-consistency.sh` never fired because committing onto an existing branch is not a feature-branch CREATE, and no gate checked which FILES land on which BRANCH). (1) **Single source of truth** — the `logic_group` → (`branch`, `track`, `path-globs`) binding is recorded ONCE in the DB and LOOKED UP, never re-invented (a commit/dispatch placing a group's file-scope on a non-canonical branch/track is a §11.4.6 no-guessing violation). (2) **Preventive hook**

(§11.4.109/§11.4.177) — a PreToolUse guard `guard-work-track-binding.sh` (inherited BY REFERENCE, never copied) BLOCKS (exit 2) a git-commit-class Bash call whose to-be-committed file set contains a path owned by a group whose canonical branch/track $\neq$ the current checkout, AND BLOCKS an Agent/Task dispatch whose §11.4.182 label track/branch $\neq$ the canonical (track, branch) of the ticket's group; fail-closed on unreadable registry, `# guardrails:allow <reason>` audited escape. (3) **Detective gate (§11.4.110/§11.4.108)** — `CM-WORK-TRACK-BINDING-ENFORCED` asserts, on the current branch's own added commits (`merge-base(HEAD,main)..HEAD`), that every `group_paths`-owned file's group destination == current branch (and `canonical_track` == current track when pinned) — the actual enforcement boundary regardless of how the commit was made, the runtime signature being FAIL-on-wrong-placement / PASS-on-canonical-placement on a clean checkout. (4) **Risk-free reconciliation (§9.2/§11.4.113/§11.4.84)** — a mis-placed change is reconciled by MERGING its commits onto the group's canonical branch (union, no commit lost, no `-s ours /reset`), ff-only push to all upstreams (NEVER force-push), then removing it from the wrong branch only after the merge is confirmed on all mirrors. Classification: universal (§11.4.17) — references no project hardware; the consuming project supplies its `(track, branch)` map, `group_paths` seed, and commit-wrapper wiring per §11.4.35. Composes

§11.4.6/.28/.29/.35/.58/.75/.84/.93/.95/.109/.110/.113/.119/.167/.176/.177/.178/.181/.182/

§9.2. Propagation gate `CM-COVENANT-114-191-PROPAGATION` (literal `11.4.191`) + recommended gate `CM-WORK-TRACK-BINDING-ENFORCED` + paired §1.1 mutation.

Canonical authority: constitution submodule `Constitution.md` §11.4.191. Non-compliance is a release blocker. No escape hatch — no `--allow-cross-track-commit`, `--skip-work-binding`, `--commit-any-branch-OK`, `--dispatch-any-track-OK`, `--no-file-scope-check` flag.

§11.4.192 — Continuous multi-track auto-backfill: a FREE track MUST be IMMEDIATELY re-assigned its next-highest-priority domain work-set, never left idle while its domain has actionable items (User mandate, 2026-07-11).

Verbatim operator mandate: "once a Track is free, we shall assign next set of workable items to it ... full continuous multi-track development ... no stopping ... this is MANDATORY". Compact summary: the moment any track's worker COMPLETES its current work-set (its item(s) closed with anti-bluff evidence, or merged), that FREE track MUST be IMMEDIATELY assigned the NEXT-highest-priority actionable work-set from its OWN domain (its §11.4.191 `logic_group` /

canonical branch) — a track is NEVER left idle while its domain still holds actionable workable items, and the conductor/ruler CONTINUOUSLY re-assigns with NO stopping (full continuous multi-track development); PROMOTES the §11.4.103(B) ≥3-parallel-background-streams-with-auto-backfill invariant into the multi-track (§11.4.176/§11.4.178/§11.4.179/§11.4.187) layer (§11.4.103(B) backfills a STREAM the moment it is done; §11.4.192 backfills a TRACK the moment its worker is done, keyed to that track's canonical (track, branch) domain per §11.4.191). (1) **WHAT** — on any track-worker completion the conductor/ruler (§11.4.187) MUST, per that track's domain, look up the NEXT actionable work-set, CLAIM it exactly-once (§11.4.176 exactly-once claim registry, no two tracks claiming one item), and DISPATCH it as the track's new §11.4.182-labelled work stream — without waiting for a scheduled wake or an operator prompt. (2) **PRIORITY** — the next work-set is the track domain's highest-priority actionable item (§11.4.42/§11.4.72 priority order, audio-first where in scope), never an arbitrary pick. (3) **SOURCE-SIDE ADVANCES EVEN WHEN VALIDATION IS GATED** — when the next item's on-device VALIDATION is blocked on a pending build / device availability (§11.4.108 clean-target / §11.4.185 manual-QA / operator-gated flash), the track MUST still advance the item's SOURCE-SIDE work (investigation, RED-test §11.4.115, source patch, pre-build gate + paired §1.1 mutation, docs) and honestly §11.4.3/§11.4.52 SKIP-with-reason ONLY the genuinely device-gated validation portion — a build/device gate on the validation step is NOT a reason to idle the whole track. (4) **HONEST BOUNDARY (§11.4.6)** — a track idles ONLY when its domain GENUINELY has zero remaining actionable workable items OR every remaining item is externally blocked with no source-side advancement possible (§11.4.94(A)/§11.4.97/§11.4.101 idle-only-when-blocked closed set); "the merges are done" / "the current item is done" / "nothing visible right now" is NEVER a valid idle reason while the domain still holds progressable items (a §11.4.94+§11.4.97 violation); an unavoidable per-item block PARKS one item, it does NOT stop the track (§11.4.101 — the track backfills its next non-blocked domain item). Classification: universal (§11.4.17) — references no project hardware; the consuming project supplies its track → domain map + priority source per §11.4.35. Composes §11.4.42/.58/.72/.94/.97/.101/.103/.167/.176/.178/.179/.182/.187. Propagation gate CM-COVENANT-114-192-PROPAGATION (literal 11.4.192) + recommended gate CM-CONTINUOUS-MULTITRACK-BACKFILL (on a track-worker completion the ruler re-assigns that track's next-highest-priority domain work-set; a track with a completed worker AND remaining actionable domain items but no new dispatch →

the gate FAILs) + paired §1.1 mutation (let a free track sit idle while its domain holds an actionable item, OR treat "merges done" as a valid idle reason → the gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.192. Non-compliance is a release blocker. No escape hatch — no `--allow-idle-track`, `--skip-backfill`, `--stop-multitrack`, `--merges-done-is-idle`, `--defer-track-reassign` flag.

§11.4.193 — Anti-blind-typing mandate: every UI interaction MUST be "seen" and "understood" via OCR/vision/screenshot proof, NEVER blind-typed (User mandate, 2026-07-13). Verbatim operator mandate: "No blind typing of any kind anywhere!!!!"; "Login to apps and services MUST BE done with understanding and seeing the whole UI ALWAYS!!!"; "blind typing is STRICTLY FORBIDDEN, and instead everything MUST BE seen and understood in coordination with proper models and means we have incorporated in HelixQA!!!!". Blind typing — sending keystrokes/text/input events to a UI surface WITHOUT first capturing AND understanding what is on screen (OCR / vision / uiautomator / accessibility-tree / screenshot) AND confirming the result after each interaction — is STRICTLY FORBIDDEN everywhere and always, for every agent/tool/automation/platform/context. It is LITERAL bluff: claiming "credentials entered"/"login completed" without visual proof the keystrokes landed on the intended fields/screen producing the intended result is a §11.4 PASS-bluff at the UI-interaction layer (a credential typed into the wrong field, or a WebView that silently swallowed input, = end-user feature broken while the agent reported success). (1) **SEE-BEFORE-YOU-TYPE** — before ANY `input text` / `input keyevent` / `adb shell input` / uiautomator gesture, capture the current screen state (screencap + Tesseract OCR ≥150 DPI, OR uiautomator dump grep, OR accessibility-tree dumphsys, OR HelixQA VisionEngine frame per §11.4.159/§11.4.160) AND confirm the intended target element is visible AND focused — the capture IS the §11.4.5/§11.4.69 `touch_input` evidence. (2) **VERIFY-AFTER-YOU-TYPE** — after EVERY interaction step, capture a NEW screen state + confirm the intended change (5 `input text` + 1 final screenshot = blind typing). (3) **OCR/VISION ORACLE MANDATORY** — Tesseract ≥150 DPI / uiautomator-dump grep / HelixQA VisionEngine / accessibility-tree dumphsys. (4) **WEBVIEW** — near-empty uiautomator hierarchy → OCR fallback per §11.4.117; if even OCR is infeasible (secure/DRM-blanked surface captures black) → the login is operator_attended-only, SKIP-with-reason per §11.4.3, NEVER blind-type-and-claim-success. (5)

CREDENTIAL-SAFE OCR — redact credential characters from OCR output + logs per §11.4.10; evidence = the redacted screenshot / OCR text. (6) **ANTI-BLUFF** — every login/interaction PASS cites a captured-evidence artefact containing pre-interaction state + post-interaction state + OCR/vision verification of the expected change; self-validated per §11.4.107(10) with a golden-good + golden-bad fixture pair. Classification: universal (§11.4.17). Composes §11.4 / §11.4.5 / §11.4.6 / §11.4.69 / §11.4.107 / §11.4.117 / §11.4.159 / §11.4.160 / §11.4.10 / §11.4.52 / §11.4.3. Propagation gate `CM-COVENANT-114-193-PROPAGATION` (literal `11.4.193`) + recommended gate `CM-NO-BLIND-TYPING` + paired §1.1 mutation (strip OCR verification before an `input text` → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.193. Non-compliance is a release blocker. No escape hatch — no `--blind-typing-ok`, `--skip-ocr-verification`, `--assume-input-landed`, `--login-without-proof`, `--type-without-seeing` flag.

§11.4.194 — Exhaustive all-scenario, all-angle code-review + verify-against-captured-runtime-evidence mandate (User mandate, 2026-07-14). Verbatim operator mandate: "CRITICAL: Code review MUST TAKE into the accounts all scenarios, look at the problem or change from all angles, deepest possible analysis! No gaps, shortcomings or false results or bluff is allowed! Extend our root constitution Submodule and all relevant files properly so the review process is more precise and we do not have situations like the last one!" STRENGTHENS §11.4.125/§11.4.134/§11.4.142 — every code-review MUST, before GO: (1) ENUMERATE the FULL input/scenario space of the change AND the defect (every dimension/path/interacting feature, not just the reported scenario); for a MULTI-FACTOR failure mode (a product of N terms, ANY degenerate term — zero/empty/null/default-init/out-of-range — triggers the failure) enumerate + INDEPENDENTLY VERIFY the fix covers EACH factor — verifying one and assuming the rest is the exact forbidden gap (forensic FACT "the last one": a boot-crash fix got a clean zero-finding GO that dismissed a clamp branch as an informational NIT "effectively unreachable ... always valid" WITHOUT PROVING it against the runtime input space; the target STILL crash-looped because the degenerate default-init input the review assumed impossible DID happen — a `size = A × B` product (e.g. `frame_size = channels × bytes_per_sample(format)`) with TWO independent zero-triggers, only ONE factor verified, the OTHER assumed away, the pre-fix crash evidence never cross-checked); (2) PROVE every assumption — "can't

happen"/"unreachable"/"always valid"/"caller never sends X" MUST be PROVEN from code+spec+captured evidence, never on faith (§11.4.6); an unproven assumption that proves false is a §11.4 review-layer bluff of PASS-bluff severity; a NIT dismissing a path as unreachable MUST carry the proof of unreachability (a bare "effectively unreachable" NIT is itself a finding); (3) VERIFY the fix against ANY captured runtime evidence (crash log/tombstone/stack-trace/sink-probe/ §11.4.108 runtime-signature/reproducer) — confirm the fix ACTUALLY prevents the CAPTURED failure against the exact captured conditions; a GO on a fix that then fails at runtime on a scenario present in the evidence (or one the review chose not to consider) is a bluff; (4) ALL ANGLES — correctness, full input domain (every factor/boundary/degenerate value), every downstream/interacting must-not-break feature (§11.4.133 target/hardware safety + shared-state features e.g. audio/multichannel/passthrough), the TESTS' OWN per-dimension coverage (a test checking only one factor of a multi-factor bug is itself a gap the review catches), the §1.1 mutation per dimension; (5) HONEST BOUNDARY (§11.4.6) — "exhaustive" = the enumerated + evidence-grounded scenario space, NOT literal omniscience; an un-analysable dimension is stated as an explicit gap (UNCONFIRMED/UNKNOWN/PENDING_FORENSICS + tracked follow-up), NEVER silently assumed safe.

Classification: universal (§11.4.17). Composes

§11.4.125/.134/.142/.6/.108/.130/.107(10). Propagation gate CM-

COVENANT-114-194-PROPAGATION (literal 11.4.194) + recommended gate CM-EXHAUSTIVE-REVIEW-ALL-SCENARIOS + paired §1.1 mutation. **Canonical**

authority: constitution submodule Constitution.md §11.4.194. Non-compliance is a release blocker. No escape hatch — no `--partial-scenario-review` , `--assume-unreachable-OK` , `--skip-captured-evidence-crosscheck` , `--single-factor-review-suffices` , `--dismiss-branch-without-proof` flag.

§11.4.195 — Branch-structure governance: taxonomy + merge-after-live-QA + flavor/product non-merge (User mandate, 2026-07-14). Verbatim operator rules: branch taxonomy `feat/<NAME_OF_FEATURE>` / `product/<NAME_OF_PRODUCT>` / `flavor/<NAME_OF_FLAVOR>` ; tracks on FEATURE branches or `main` MUST — after full confirmation + validation + verification + full LIVE MANUAL TESTING (§11.4.185) — MERGE to `main` COMPLETELY (all submodules + main repo), ff-only, NEVER force-push (§11.4.113), no commit lost (§9.2); tracks with an entirely-new FLAVOR/PRODUCT do ALL work on a `flavor/` AND/OR `product/` branch (both may exist) which does NOT merge to `main` the same way; commit + push ALL branches to ALL upstreams (§2.1), ff-only. Every controlled branch belongs to

one of THREE canonical taxonomy classes, each with a distinct integration lifecycle: **(A) Taxonomy** (extends §11.4.181/§11.4.167(D)) — `feat/<slug>` = FEATURE (merges to `main` after live-manual-QA, clause B); `product/<slug>` = PRODUCT line + `flavor/<slug>` = FLAVOR line (do NOT merge to `main` the standard way, clause C); `<slug>` lowercase snake/kebab (§11.4.29); the canonical branch NAME is minted ONCE + used IDENTICALLY across main repo + every touched owned submodule (§11.4.181), recorded once in the §11.4.176 claim registry / §11.4.93 `logic_group` → `branch` map + LOOKED UP never re-invented (§11.4.6); the prefix STRINGS are operator-canonical, a project already on a different feature-prefix (`feature/...`) is compliant so long as it is the ONE canonical name per §11.4.181, reconciled by the §11.4.181 risk-free MERGE-onto-canonical path (ff-only, no commit lost, never a rename that drops commits). **(B) FEATURE/main** → **merge-after-live-QA** — before merging to `main` satisfy IN ORDER automated four-layer (§11.4.4(b)) + §11.4.40 full-suite retest THEN full LIVE MANUAL TESTING by the QA team (§11.4.185 FINAL human gate, automated GREEN necessary NOT sufficient); only after §11.4.185 may the merge proceed + it MUST be COMPLETE (main repo + every touched owned submodule) ff-only onto latest `main` (§11.4.113), NEVER force-push, NO commit lost / NO history rewrite (§9.2/§11.4.41 step 3), pushed to ALL upstreams (§2.1); binds the §11.4.185 gate as an explicit precondition of the §11.4.167(I)/§11.4.188 approval-gated back-merge. **(C) FLAVOR/PRODUCT lines** → **do-not-merge-to-main-the-same-way** — an entirely-new flavor/product line does ALL work on a `product/<NAME>` AND/OR `flavor/<NAME>` branch (both may coexist), a divergent line with its own release lifecycle that does NOT flow back to `main` via clause B; STILL obeys §11.4.113 (no force-push) / §9.2 (no commit lost) / §2.1 (push all branches all upstreams ff-only) / §11.4.188 (MAY merge `main` INTO itself to stay close — trunk-into-line permitted; only line-into-trunk is gated/withheld). **(D) Push discipline** (§2.1/§11.4.113) — commit + push ALL branches (feature/product/flavor) to ALL upstreams ff-only, NEVER force-push; new-branch creation is inherently ff (§9.2) + non-disruptive from an existing HEAD, permitted in a dirty tree under §11.4.84 provided no commit is in-flight. Honest boundary §11.4.6 — governs branch STRUCTURE + integration DIRECTION, does NOT prove any branch's work correct (still §11.4.108 runtime-signature + §11.4.40 retest + §11.4.185 manual-QA). Classification: universal (§11.4.17) — strict generalisation of §11.4.181 + §11.4.167(D)/(I) + §11.4.188 + §11.4.185. Composes §2.1/§9/§9.2/§11.4.6/§11.4.17/§11.4.29/§11.4.35/§11.4.40/§11.4.41/§11.4.84/§11.4.93/§11.4.113/

§11.4.167(D)/(I)/§11.4.176/§11.4.178/§11.4.181/§11.4.185/§11.4.188/§11.4.191.
 Propagation gate `CM-COVENANT-114-195-PROPAGATION` (literal `11.4.195`) +
 recommended gate `CM-BRANCH-TAXONOMY-CLASS` (every controlled branch prefix
 $\in$ `feat/` | `product/` | `flavor/` | reserved `main` / `master` ; a `product/` |
`flavor/` branch is NOT the source of a merge onto `main` ; a `feat/` | `main`
 back-merge cites a §11.4.185 manual-QA confirmation) + paired §1.1 mutation (a
`product/*` merged onto `main` , OR a `feat/*` back-merge with no §11.4.185
 citation → gate FAILs; strip the literal → propagation gate FAILs; gate-code =
 separate work item). **Canonical authority:** constitution submodule
`Constitution.md` §11.4.195. Non-compliance is a release blocker. No escape
 hatch — no `--any-branch-prefix-OK` , `--merge-product-to-main` , `--skip-`
`live-qa-before-merge` , `--force-push-branch` ,
`--partial-submodule-merge` flag.

§11.4.196 — Native-alias-first priority + per-alias real-signal limit/subscription tracking + auto-rebind-on-recovery + resource-detection-by-real-identity-not-substring-match (User mandate, 2026-07-14). Verbatim operator mandate: use ALL operational claude-NATIVE aliases FIRST — as long as ≥ 1 native alias is operational (NOT session-rate-limited, NOT weekly-limit-reached, NOT expired-subscription) — before ANY provider alias; track per alias the session rate-limit / weekly-limit / subscription expiry+renewal from REAL signals (never faked, §11.4.6) + record when each becomes operational AGAIN; the moment a higher-priority native recovers, rebind tracks to it; and fix the host-budget guard footgun where a `pgrep -f substring-match false-REFUSED` every worker spawn. The multi-track alias-binding + rate-limit/subscription-tracking layer of the §11.4.187 ruler orchestration MUST implement ALL of: **(A) NATIVE-ALIAS-FIRST PRIORITY** — every OPERATIONAL native (`claudeN`) alias is selected before ANY provider, UNCONDITIONALLY, while ≥ 1 native is operational; the guarantee is STRUCTURAL — a native/provider CLASS partition scanned native-class-completely-before-provider-class (§11.4.111 resolve-by-CLASS-not-by-file-position), so a provider NEVER wins while a non-cooled native exists regardless of roster order; within a class the operator-mandated order (native equal-capability, then providers `deepseek` → `xiaomi` → `opencode` → `kimi` → ...) is authoritative decision-DATA (§11.4.28) never re-invented (§11.4.6). **(B) PER-ALIAS REAL-SIGNAL LIMIT + SUBSCRIPTION TRACKING** — "operational?" per alias from REAL captured signals, NEVER faked/guessed (§11.4.6): the §11.4.187 captured 429 signature (`isApiErrorMessage` + `apiErrorStatus:429` / `api_error_status:429`) sub-

classified into reason-CLASS closed set {session | weekly | subscription} from real markers (reset/weekly-limit/session-limit); each limited alias records its CLASS + an operational-again EPOCH with a DIFFERENT window per class — session (self-heals at reset, short), weekly (~until the weekly reset, long), subscription (INDEFINITE — operational-again UNKNOWN until a REAL renewal signal, parked with a far-future sentinel; §11.4.6 NEVER guesses a renewal date, an operator/config supplies it). A limited alias is NEVER auto-selected until its epoch passes (session/weekly) or a real renewal/recovery clears it (subscription).

(C) AUTO-USE-ON-RECOVERY — the moment a higher-priority native recovers (window elapsed / renewal marked), rebind tracks UPWARD to it (a promote -class op keyed off a comparable priority-rank: native class before provider, then config order), PRESERVING each track's worktree AND device leases (leases key on the STABLE track id — the alias changes, not the track; §11.4.119/§11.4.187); idempotent. **(D) RESOURCE-DETECTION BY REAL IDENTITY, NOT SUBSTRING-MATCH (RB-02 guard fix)** — a "heavy work in flight" guard MUST identify a real resource by its ACTUAL process command line read from the OS (/proc/<pid>/cmdline), NEVER a bare pgrep -f REGEX substring match that false-matches a process merely MENTIONING a token — a pattern-CARRIER (a process quoting the whole alternation pattern; a claude / worker whose prompt embeds the token; a launcher/monitor/test), a same-tool ENGINE process, or the guard's own process/parent (the §12.12 self/sibling pgrep footgun — forensic FACT 2026-07-13: the guard false-REFUSED EVERY §11.4.187 spawn on a host with ZERO real builds because a live claude worker's prompt quoted the pgrep pattern); the guard re-reads each matched PID's real cmdline + excludes those carrier/engine/worker/self classes; an UNREADABLE/ vanished cmdline is conservatively counted as the real resource (REFUSE — safe reversible default, §11.4.101/§12.8), so a genuine build is still caught. **(E) ANTI-BLUFF (§11.4/§11.4.108)** — four-layer coverage each (pre-build gate + on-host test + paired §1.1 mutation + captured evidence §11.4.69) with §11.4.115 RED-polarity + §11.4.50 determinism; honest boundary §11.4.6 — mandates the MECHANISM + REAL-signal derivation of "operational", genuine live per-subscription quota-isolation proof still needs live tokens (operator-scheduled acceptance window), never conflated with a mock/hermetic-contract GREEN. Classification: universal (§11.4.17). Composes §11.4.187/.182/.176/.111/.101/.103/.6/.108/.115/.119/§12.12/§12.6/§12.8. Propagation gate CM-COVENANT-114-196-PROPAGATION (literal 11.4.196) + recommended gate CM-ALIAS-PRIORITY-LIMIT-TRACKING (native-first CLASS

partition; per-alias reason-class + operational-again tracking; auto-rebind-on-recovery; build-guard resolves by real cmdline not bare substring) + paired §1.1 mutation (swap native/provider priority bases so a provider outranks a native → recommended gate FAILs; flip the guard's strict-cmdline-filter to the pre-fix bare `pgrep` → its RB-02 test FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.196. Non-compliance is a release blocker. No escape hatch — no `--provider-before-operational-native`, `--fake-limit-state`, `--guess-renewal-date`, `--no-rebind-on-recovery`, `--substring-match-build-guard`, `--skip-real-signal-classification` flag.

§11.4.197 — Research / kicked-off-work completion mandate: every research effort MUST be driven to full, wired, verified completion or explicitly closed — never left un-wired in the backlog (User mandate, 2026-07-15). Verbatim operator mandate: "all research work we start MUST BE fully finished, completely confirmed, validated and verified! It MUST NOT happen for it to stay at the backlog uncompleted and un-wired! Loss of requirements, failure to incorporate them is FORBIDDEN!!! Everything MUST BE fully completed! No exceptions!" Every research effort / kicked-off improvement / design doc / spike / prototype / PoC / partially-landed feature a project STARTS MUST be driven to FULL completion OR explicitly evidence-backed CLOSED — there is no third "sitting un-wired in the backlog" state; "we researched/designed/started it" is NEVER terminal, and a `docs/research/*` doc / design doc / spike / half-wired feature that is neither COMPLETED-and-wired nor explicitly-closed is a §11.4.197 violation of §11.4 PASS-bluff severity at the requirements-integrity layer (the requirement was accepted, work began, then it silently evaporated); loss of requirements + failure to incorporate started work are FORBIDDEN. **COMPLETED** = (1) implemented + (2) WIRED (integrated AND used by default, NOT merely present behind a dead flag / uncalled function / unreferenced module / never-selected path — mere-presence-without-wiring is the §11.4.124 dead-code + §11.4.108 layer-2/3 SOURCE → ARTIFACT → RUNTIME gap) + (3) confirmed/validated/verified with captured physical evidence §11.4.5/§11.4.69/§11.4.107 on a clean deployment §11.4.108 runtime-signature (never metadata-only/config-only/absence-of-error/grep-without-runtime §11.4/§11.4.1) + (4) tracked to a terminal Status §11.4.15/§11.4.33. **Explicitly CLOSED** = deliberately terminated with a documented evidence-backed reason from the closed vocabulary — §11.4.90 Obsolete / §11.4.112 structurally-impossible / §11.4.122 operator-confirmed drop — NEVER

silent abandonment ("no time"/"lower priority" = §11.4.21 Operator-blocked or a still-open tracked item, NOT a close). **Mechanical tracker binding (§11.4.93):** every research effort / design doc / kicked-off improvement MUST map to ≥1 tracked workable item (ATM-NNN §11.4.54) whose completion the loop enforces; a research doc with NO implementing item, or a tracked item stalled un-wired past the cadence, is a violation; the autonomous loop (§11.4.87/.94/.97/.103/.126/.192) MUST treat every un-wired research effort as an actionable queue item — driving it to COMPLETED or explicitly CLOSED is exactly the "keep working until the queue is genuinely empty" contract. Honest boundary §11.4.6 — "fully completed" is bounded by wired+verified OR evidence-backed-closed, NOT "every speculative idea ships"; every STARTED effort reaches a documented terminal state, proven never assumed. Classification: universal (§11.4.17). Composes §11.4.8/.87/.94/.97/.103/.118/.123/.124/.150/.192/.90/.112/.122/.21/.93/.54/.108/.15/.33/.126. Propagation gate `CM-COVENANT-114-197-PROPAGATION` (literal `11.4.197`) + recommended gate `CM-RESEARCH-COMPLETION-TRACKED` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.197. Non-compliance is a release blocker. No escape hatch — no `--research-may-stay-unwired`, `--skip-completion-tracking`, `--abandon-without-close`, `--present-but-unwired-is-done`, `--backlog-research-OK`, `--lose-requirement-OK` flag.

§11.4.198 — Default working mechanisms: multi-alias native-first orchestration AND heavy token-optimization are ALWAYS-ON defaults for single- and multi-track work (User mandate, 2026-07-15). Operator mandate: multiple claude-toolkit aliases (native-first) AND heavy token-optimization MUST be the DEFAULT mechanisms, always used with single- and multi-track development and with single or multiple working agents. Two proven mechanisms promoted to STANDING always-on defaults, engaged automatically for EVERY configuration (single- OR multiple-agent, single- OR multi-track), never a per-request opt-in: **(A) MULTI-ALIAS NATIVE-FIRST ORCHESTRATION** — the claude-toolkit multi-alias mechanism (every OPERATIONAL claude-native alias before ANY provider per §11.4.196 native-alias-first priority + per-alias real-signal limit/subscription tracking + auto-rebind-on-recovery, driven by the §11.4.187 automatic multi-track ruler) MUST be the default execution substrate for ALL work, NOT only the multi-track case; a SINGLE-track/SINGLE-agent session STILL binds its worker through the same native-first alias-priority mechanism (best operational native + rebind-on-recovery), a MULTI-track/multiple-agent config uses it across every track (§11.4.176

exactly-once claim + §11.4.182 track+branch+alias label); running on an arbitrary/hardcoded/provider-first alias when an operational native exists, or bypassing the mechanism because "it is just one agent", is a violation. **(B) HEAVY TOKEN-OPTIMIZATION** — the §11.4.141 token-efficiency discipline (compact-anchor layout with full-one-hop-away, targeted reads over whole-file dumps, subagent context isolation §11.4.20/§11.4.70, evidence-hash references not evidence dumps) **MUST** apply **ALWAYS** — every read/dispatch/doc-edit — for single- AND multi-track, single OR multiple agents; token-heavy patterns (re-reading whole large files when a targeted slice suffices, dumping full evidence into context, monolithic non-subagent execution of parallelisable work) when a token-efficient path exists are violations. Honest boundary §11.4.6 — "default" = auto-engaged without an operator request; it does NOT weaken host-safety, the multi-alias/multi-agent fan-out stays under the §12.6 60% memory ceiling + §12.12 thread-headroom + §11.4.58 parallel-agent cap (contending/redundant agents are NOT "the default", §11.4.183 maximum-USEFUL-throughput); §11.4.198 BINDS the two mechanisms as always-on, it does NOT redefine them (mechanics live in §11.4.196/§11.4.187/§11.4.141). Classification: universal (§11.4.17). Composes §11.4.141/.187/.196/.182/.176/.183/.103/.58/.20/.70/.126/§12.6/§12.12. Propagation gate `CM-COVENANT-114-198-PROPAGATION` (literal `11.4.198`) + recommended gate `CM-DEFAULT-MECHANISMS-ALWAYS-ON` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.198. Non-compliance is a release blocker. No escape hatch — no `--single-agent-skips-alias-priority`, `--provider-first-default`, `--token-optimization-optional`, `--heavy-tokens-OK`, `--not-default-mechanism` flag.

§11.4.199 — Exact-reproduction-sequence mandate: when a working reproduction exists, the investigation MUST use ITS exact sequence — a deviating repro proves NOTHING (research-derived, 2026-07-15). Forensic FACT (2026-07-15): an investigation concluded a defensive "cover" mechanism "never engages" (and the defect it guards was therefore unfixed) — the conclusion was **WRONG** because the hand-rolled reproduction drove the target with a **DIFFERENT** input than the existing test (a HOME key where the test sent a media-pause key), so it never reached the PRECONDITION the mechanism keys off; re-driven with the TEST'S EXACT sequence the mechanism provably **ENGAGED**. Every investigation / debugging pass / validation that reproduces a defect or exercises a code path **MUST**: (1) **use the existing reproduction's EXACT sequence** when a working reproduction already exists (a test/harness/recorded

sequence that demonstrably reaches the defect) — same inputs, same order, same timing, same preconditions, same topology (§11.4.3); a hand-rolled approximation is NOT a substitute; (2) **a deviating variant MUST FIRST prove it reaches the precondition** with captured evidence before ANY conclusion is drawn from it; (3) **a repro that misses the precondition proves NOTHING** — it is NOT evidence of absence (§11.4.6), and concluding "defect present"/"defect absent"/"the mechanism never engages" from it is a §11.4/§11.4.1 bluff at the INVESTIGATION layer (a FAIL-bluff when it wrongly condemns working code, a PASS-bluff when it wrongly clears broken code); (4) **record the exact driving sequence as captured evidence** so it is replayable + auditable, every deviation explicit + justified + precondition-proven. STRENGTHENS §11.4.102 (Phase-2 reproduce) / §11.4.146 (reproduce-first) / §11.4.115 (RED on the broken artifact) — they mandate THAT you reproduce, §11.4.199 mandates HOW: with the sequence PROVEN to reach the defect. Classification: universal (§11.4.17). Composes §11.4.1/.3/.5/.6/.7/.102/.107/.115/.123/.135/.146. Propagation gate CM-COVENANT-114-199-PROPAGATION (literal 11.4.199) + recommended gate CM-EXACT-REPRO-SEQUENCE + paired §1.1 mutation. **Canonical authority:** constitution submodule Constitution.md §11.4.199. Non-compliance is a release blocker. No escape hatch — no --approximate-repro-OK, --skip-precondition-proof, --hand-rolled-repro-suffices, --absence-proves-absence, --deviating-sequence-conclusion-OK flag.

§11.4.200 — Non-targetable deploy/flash tooling MUST isolate exactly ONE eligible target and VERIFY-AFTER-WRITE on the INTENDED target (research-derived, 2026-07-15). Forensic FACT (2026-07-15): a flashing tool with no device-selection argument AUTO-SELECTED a device from the bus, silently wrote the freshly-built image to the WRONG board (a sibling still attached), printed SUCCESS, and the INTENDED board kept its OLD build — every downstream test then validated a target that had NEVER received the artifact (the §11.4.108 ARTIFACT → RUNTIME layer silently broken while every tool reported green). Whenever a deploy / flash / install / write tool CANNOT address a specific target, the procedure MUST: (1) **targetability is a PRECONDITION** — EITHER pass an explicit STABLE selector (serial/UUID/stable id per §11.4.111, never an enumeration ordinal) OR GUARANTEE exactly ONE eligible target is present at the moment of the write (siblings disconnected/powered-down/excluded), the enumerated eligible-target set captured as evidence (size == 1); (2) **the tool's own success message is NOT proof** — an "ok"/exit-0 proves only that bytes reached

the device the TOOL selected, NEVER the device the OPERATOR intended (§11.4.6); (3) **VERIFY-AFTER-WRITE on the intended target** — read back the deployed artifact's IDENTITY (build id / version / checksum / fingerprint) FROM the intended target and assert it equals the artifact just written; a mismatch or the OLD identity is a FAIL, never a warning (the §11.4.108 ARTIFACT → RUNTIME assertion applied to the deployment step); (4) **no validation on an unverified target** — any test run against a target whose deployed-artifact identity was NOT verified is INVALID and any PASS is a §11.4 PASS-bluff. Honest boundary §11.4.6 — isolation + read-back prove the INTENDED target holds the INTENDED artifact, NOT that the artifact is correct (§11.4.108/.40/.130). Classification: universal (§11.4.17). Composes §11.4.6/.46/.108/.111/.119/.130/.133/.139. Propagation gate `CM-COVENANT-114-200-PROPAGATION` (literal `11.4.200`) + recommended gate `CM-DEPLOY-TARGET-ISOLATED-AND-VERIFIED` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.200. Non-compliance is a release blocker. No escape hatch — no `--auto-select-target-OK`, `--trust-tool-success-message`, `--skip-post-write-verify`, `--validate-unverified-target`, `--multiple-targets-on-bus-OK` flag.

§11.4.201 — Every guard/gate MUST assert the REAL condition: a false-positive refusal is a FAIL-bluff, a false-negative pass is a PASS-bluff (research-derived, 2026-07-15). Forensic FACT (2026-07-15, two independent instances in one cycle): (a) a host-budget guard used a bare `pgrep -f '<pattern>'` and REFUSED every worker spawn on a host with ZERO real builds because a live worker's own command line merely QUOTED the pattern (the §11.4.196(D)/§12.12 carrier footgun); (b) a host-safety pre-flight refused a build on a "swap full" reading that did not reflect a real memory-pressure condition, clearing only after the operator manually cycled swap — BOTH guards refused on a PROXY signal that was NOT the real condition. Every GUARD / GATE / PRE-FLIGHT (host-safety, resource-budget, readiness, lock, capability, availability) MUST assert the REAL condition it CLAIMS, resolved from the AUTHORITATIVE source (the real `/proc/<pid>/cmdline`, the real cgroup/meminfo pressure metric, the real lock-holder liveness via `kill -0`, the real device identity), NEVER from a proxy signal that something which is NOT the condition can satisfy: (1) **a FALSE-POSITIVE REFUSAL is a FAIL-bluff (§11.4.1)** — blocking work when the condition is ABSENT is exactly as forbidden as passing while a defect is present (it halts real work AND teaches operators/agents to bypass guards); (2) **a FALSE-NEGATIVE PASS is a PASS-bluff (§11.4);** (3) **SELF-VALIDATED (§11.4.107(10))** — every

guard ships a golden-TRUE fixture (condition present → guard FIRES) AND a golden-FALSE fixture INCLUDING the known CARRIER/decoy case (condition absent but proxy signal present → guard MUST NOT fire), both wired into the meta-test; a guard firing on its golden-FALSE fixture is ITSELF the defect; (4)

conservative-safe default on an unresolvable signal (§11.4.101/§12) — refuse the destructive/high-blast-radius action and say so HONESTLY ("could not resolve the condition" is NOT a false positive; claiming a condition never verified IS one, §11.4.6); (5) **every refusal reports its RESOLVED evidence** (real pid+cmdline, real metric+threshold, real lock holder) so a false positive is diagnosable in one step. GENERALISES §11.4.196(D) (resource-detection-by-real-identity — to EVERY guard) + §11.4.180 (liveness PROVEN before reaping — to every guard's condition). Classification: universal (§11.4.17). Composes §11.4.1/.6/.101/.107(10)/.109/.111/.180/.196(D)/§12/§12.6/§12.8/§12.12.

Propagation gate `CM-COVENANT-114-201-PROPAGATION` (literal `11.4.201`) + recommended gate `CM-GUARD-ASSERTS-REAL-CONDITION` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.201. Non-compliance is a release blocker. No escape hatch — no `--proxy-signal-OK`, `--substring-match-guard`, `--false-positive-refusal-acceptable`, `--unvalidated-guard`, `--refuse-without-evidence` flag.

EXTENSIONS landed 2026-07-15 (existing anchors strengthened — full text in `Constitution.md`). §11.4.89(G) subagent long-op discipline + conductor-owns-inherently-long-ops: a SUBAGENT MUST NOT run a single synchronous op exceeding the runtime's no-progress stream watchdog (~5 min) — it backgrounds + polls (each poll IS the progress the watchdog observes); inherently-long ops (full builds / full suites / device flashes) are the CONDUCTOR's to run (no such watchdog); a stall-killed subagent is a CRASH per §11.4.147 (respawned, never absorbed as a completion); watchdog budget read from the runtime, never guessed (§11.4.6). **§11.4.106(F)** write-seam hook enforcement: the doc/DB sync MUST be enforced at the COMMIT seam (hook refuses a commit whose staged set touches a chain source while its exports/DB are stale) + the BUILD seam (pre-build gate) + the CONSTITUTION-PULL seam (§11.4.164 post-update hook) — a stale CONTINUATION (§12.10) / session-resumption file (§11.4.131) / memory artefact while live state moved is a violation; honest boundary §11.4.6 — seam hooks give eventual-consistency-at-every-write, NOT literal real-time (that needs a watch daemon, carried as a tracked §11.4.197 upgrade path, never claimed as shipped). **§11.4.147(e)** API-QUOTA / RATE-LIMIT EXHAUSTION is a FIRST-CLASS CRASH

CLASS: an agent killed mid-work by a weekly/session/subscription quota boundary is `crashed` (never `complete`) — record the REAL limit signal + reason-class per §11.4.196(B), REBIND to the next OPERATIONAL alias per §11.4.196(A)/(C) and RESPAWN there (bounded-PARK per §11.4.101 only when EVERY alias is cooled), RESUME from the killed agent's captured transcript + partial state at its EXACT last point, and the §11.4.87/.94/.97/.126 loop done-condition MUST NOT read satisfied while any quota-killed entry is non-`complete`. **§11.4.196(F) CONFIGURED ≠ IN USE: forensic FACT** — the alias orchestrator was installed with a valid config yet its binding registry was EMPTY while THREE live workers ran outside it, so `promote` rebound ZERO tracks; EVERY spawn path MUST bind THROUGH the orchestrator, and the mechanism's readiness is its RUNTIME SIGNATURE (§11.4.108) — **the binding-registry rows MUST EQUAL the live-worker set** — an empty/partial registry alongside live workers is a FALSE-READY state and claiming "ready/in use" on config alone is a §11.4/§11.4.6 orchestration-layer bluff.

§11.4.142 + §11.4.194 re-affirmed (2026-07-15 audit, no change required): the operator re-insisted 3× this session that EVERY change passes an exhaustive code review, ALWAYS, before build/commit. Audited: §11.4.142 already binds EVERY change to ANY governed repository — "no exception ... no 'just a doc edit' exemption, no 'just a one-liner' exemption, no §11.4.92-self-review substitution, no 'trivial change' carve-out ... BEFORE it is accepted, committed, or built" — with independence load-bearing (§11.4.70/§11.4.20); §11.4.194 supplies the exhaustive all-scenario / prove-every-assumption / cross-check-captured-runtime-evidence precision bar; §11.4.134 iterates the review to a ZERO-finding, ZERO-warning GO; §11.4.125 places the gate before the pre-build sweep + main build. NO carve-out exists in any of them — the family is complete as written and needed no widening this round (§11.4.6: audited, not assumed).

§11.4.202 — Reporting directives: a report MUST auto-create a fully-populated, fully-synced workable item — never a prose acknowledgement (User mandate, 2026-07-15). Verbatim operator mandate: introduce reporting directives `ISSUE` / `BUG` / `TASK` that create a proper workable item filled with all details, with the WHOLE workable-items system automatically up to date and fully in sync — the main DB, all documents, and all external tracking systems; universal, reusable on every project, fully decoupled, recognized immediately by every project that pulls the constitution. Every governed project MUST expose REPORTING DIRECTIVES — the §11.4.140 registered actions `BUG` (Type=Bug), `TASK` (Type=Task), and `ISSUE` (generic report, CLASSIFIED into the §11.4.16 closed

set) — such that a plain-language report from the operator is TURNED INTO a real, tracked, fully-synced workable item, automatically, every time. A report that is discussed, acknowledged, answered in prose, or "noted" WITHOUT landing a tracked item is a §11.4.202 violation of §11.4 PASS-bluff severity at the requirements-intake layer: the requirement was accepted and then silently evaporated (the §11.4.197 loss-of-requirements failure, at its entry point). **(1) GRAMMAR (§11.4.140).** The three directives are ordinary registry rows, so EVERY §11.4.140 form works out of the box (`BUG :: <text>` , `DEFAULT::BUG :: <text>` , `/DEFAULT::BUG <text>` , `BUG ---> <text>`), PLUS a SIXTH form this anchor adds — the **single-colon form** `NAME: <text>` (`ISSUE: subtitles are late`), the shape operators actually type. The single-colon form is **REGISTERED-ACTION-ONLY by construction**: a single-colon token that is NOT a registered action is a NO-OP (an ordinary prompt), NEVER an ASK and NEVER a sub-system shortcut — because ordinary English prose (`TODO:` , `WARNING:` , `FIXME:`) shares that shape (`NOTE:` is now a registered §11.4.140 severity marker), and routing it to the §11.4.66 clarify path would question every such sentence (§11.4.6 honest boundary: the other five forms keep their ASK-on-unknown behaviour — their shapes do not occur in prose). Host-command collisions are handled by the EXISTING §11.4.140 conflict mechanism (`slash_bare: auto + slash_conflicts: [...]`), which already satisfies the operator's explicit-prefix requirement: a bare `/NAME` that collides with a built-in (e.g. the documented Claude Code built-in `/bug`) is NOT honored, while the namespaced `/DEFAULT::NAME` form is ALWAYS unambiguous — and the `NAME: / NAME :: / NAME --->` forms are never affected by any host command.

(2) THE ITEM. The created item MUST carry, on ALL surfaces: a valid Status (§11.4.15, `Queued` at intake) + a Type from the CLOSED set {Bug | Feature | Task} (§11.4.16 — `ISSUE` is a REPORTING CHANNEL, never a fourth type; inventing one is a violation) + a stable auto-incremented id (§11.4.54), and a COMPREHENSIVE structured description (§11.4.148 D2 / §11.4.171): what / affected-scope / reproduction / acceptance. An undetermined section is recorded as an explicit `UNKNOWN:` gap — NEVER invented (§11.4.6). For `ISSUE` , the type is classified from the report's content and stated as FACT; if the content does not determine it, the operator is ASKED (§11.4.66 / §11.4.105) BEFORE creation; ONLY when running autonomously where asking is impossible (§11.4.101) may the §11.4.16 lowest-stakes ambiguity default `Task` be used — and then the defaulted classification MUST be recorded verbatim in the item and surfaced for

reclassification, never silently asserted. **(3) THE FULL SYNC (the load-bearing half).** Creating the item is NOT sufficient: the directive MUST drive the WHOLE workable-items system into sync, in one shot — (a) the item lands in the SQLite single-source-of-truth (§11.4.93 / §11.4.95); (b) EVERY derived document is regenerated FROM the DB — the open + closed trackers, their summaries, and their `.md` / `.html` / `.pdf` / `.docx` siblings (§11.4.12 / §11.4.53 / §11.4.65 / §11.4.106); (c) the item is pushed to EVERY configured external tracker (§11.4.148 D5). **(4) NEVER A FAKED PUSH (§11.4.10 / §11.4.6 / §11.4.3).** A tracker whose credentials or whose client are absent is SKIPPED with an honest machine-readable reason (`credentials_absent` — names of the unset variables ONLY, never a value; `tracker_client_absent` — PENDING-OPERATOR-INPUT). A tracker PASS may be recorded ONLY after a real push command really exited 0. Fabricating, assuming, or silently omitting a push is a §11.4 bluff at the integration layer; a missing tracker NEVER blocks the item from landing in the DB + docs (the report is never lost because a mirror is unreachable). **(5) DECOUPLED ENGINE, CONSUMER-OWNED DATA (§11.4.28 / §11.4.177).** The item-creation + full-sync engine lives in the constitution submodule (`constitution/scripts/reporting/report_item.sh`), is inherited BY REFERENCE (never copied), and carries ZERO project literals; every project-specific value — DB path, id prefix, sync command, tracker commands + their required env vars — comes from a CONSUMER-OWNED config file (DATA, never an engine edit). The engine FAILS CLOSED with an actionable message when that config is absent — it never guesses a project's paths (§11.4.6). **(6) ANTI-BLUFF (§11.4 / §11.4.107(10)).** The mechanism ships four-layer coverage: a pre-build gate (`CM-REPORTING-DIRECTIVES`) whose FUNCTIONAL invariant SOURCES and RUNS the parser (never a grep-only assertion, §11.4.108), a paired §1.1 mutation test that proves each invariant is genuinely load-bearing, and an end-to-end suite run against a REAL SQLite DB with REAL binaries (no mocks, §11.4.27) that asserts a real row lands with the right Type/Status/id, that the docs are regenerated FROM the DB, and that an absent tracker SKIPS honestly — self-validated with a golden-good + golden-bad + negative-control fixture set (an oracle that passes its golden-bad fixture is itself the bluff). Classification: universal (§11.4.17) — the consuming project supplies its DB path, id prefix, sync command, and tracker bindings per §11.4.35. Composes §11.4.140 (grammar) / §11.4.93 / §11.4.95 (DB SSoT) / §11.4.15 / §11.4.16 / §11.4.54 (status + type + id) / §11.4.148 / §11.4.171 (comprehensive description) / §11.4.106 / §11.4.12 / §11.4.53 / §11.4.65 (doc sync) / §11.4.10 (credentials) / §11.4.66 / §11.4.105 (clarify, never guess) / §11.4.101 (autonomous decision) /

§11.4.197 (a started requirement never evaporates) / §11.4.28 / §11.4.177 (decoupling) / §11.4.164 (auto-registration on constitution pull) / §11.4.107(10) / §1.1. Propagation gate `CM-COVENANT-114-202-PROPAGATION` (literal `11.4.202`) + recommended gate `CM-REPORTING-DIRECTIVES` (registry declares ISSUE/BUG/TASK + the single-colon form is registered-only; the engine exists, is parse-clean, and is project-literal-free; the three directives really expand at runtime while prose does not; the honest-skip reasons exist and a tracker PASS is gated on a real exit 0) + paired §1.1 mutation (strip an action row, flip the single-colon form to non-registered-only, strip the parser branch, remove the prose NO-OP branch, make the engine fake a tracker push, or couple the engine to one project → the gate FAILs; strip the literal → the propagation gate FAILs).

Canonical authority: constitution submodule `Constitution.md` §11.4.202.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--report-without-item`, `--prose-acknowledgement-OK`, `--skip-item-creation`, `--skip-tracker-sync`, `--fake-tracker-push`, `--invent-a-fourth-type`, `--hardcode-project-in-engine` flag.

§11.4.207 — Instant multi-stream resume engine: whole-fleet continuation state is a durable, content-addressed, atomically-committed snapshot resumed in $O(\text{changed})$ reads, not an $O(\text{history})$ re-read (User mandate, 2026-07-15). Compact summary: the continuation mechanism (§12.10 / §11.4.127 / §11.4.131 / §11.4.205) is EXTENDED — never replaced — by a mechanical engine so a fresh session resumes ALL work (every Track, main stream, and agent) at once, INSTANTLY and at token cost $O(\text{changed streams})$, by reading ONE snapshot manifest + one compact blob per stream, NEVER re-reading the full handoff history. (1) **Content-addressed Merkle store** — each stream's canonical state bytes are named by their sha256 (dedup: unchanged streams share a hash → a global fleet snapshot costs $O(\text{changed})$); a snapshot is ONE manifest referencing every stream's content-hash (§11.4.205(4)). (2) **$O(\text{streams})$ resume, $O(\text{history})$ NEVER on the resume path** — `RestoreAll` reads the manifest + one blob per stream; the append-only event ledger is audit + lost-update recovery only, never read to resume (§11.4.127). (3) **Tamper-evident + deterministic** — every blob is integrity-checked on read (bytes MUST hash to their id) and a deterministic round-trip (re-serialize must byte-match), no wall-clock in the hashed content, verified by a SELF-VALIDATING oracle: golden-good PASS + golden-bad detected FAIL + negative-control (a legitimately-older-but-valid snapshot) PASS — the false-

positive guard proving the verifier distinguishes a LAGGING copy from a TAMPERED one (§11.4.201/§11.4.107(10)/§11.4.206(3)). (4) **Durable + atomic + single-writer + crash-safe** — atomic ref write (temp → fsync → rename → dir-fsync, §11.4.205(6)), append-only ledger (§11.4.205(6)), single-writer-per-stream refusal (§11.4.206), advisory lock with provably-stale reap (`kill -0` liveness, never steals a live lock, §11.4.180/§9.2). (5) **Project-agnostic engine, consumer-owned data** — the engine (`github.com/vasic-digital/continuum` , GitLab mirror) carries ZERO project literals, fails closed rather than guess a store path (§11.4.6/§11.4.177), and is inherited BY REFERENCE as a depth-1 submodule under the §11.4.28(C) carve-out (`constitution/submodules/continuum/` , `helix-deps.yaml` , zero own-org deps); every consumer supplies its store path / actor / stream naming as DATA. Anti-bluff (§11.4/§11.4.108): the whole engine ships four-layer coverage (unit + integration + e2e + stress + chaos §11.4.85, race-clean) with a measured resume Metric (`blobs_read == 1 + streams` is the O(streams) proof) and the self-validating oracle as captured evidence (§11.4.5/§11.4.69/§11.4.107); honest boundary (§11.4.6): it snapshots the state an agent REPORTS, not its execution image (not a CRIU/durable-execution runtime, §11.4.112), and `est_tokens` is a documented `bytes/4` heuristic, not a tokenizer count. Classification: universal (§11.4.17). Composes §12.10 / §11.4.127 / §11.4.131 / §11.4.205 / §11.4.116 / §11.4.147 / §11.4.180 / §11.4.201 / §11.4.206 / §11.4.107(10) / §11.4.28 / §11.4.177 / §11.4.85 / §9.2. Propagation gate `CM-COVENANT-114-207-PROPAGATION` (literal `11.4.207`) + recommended gate `CM-CONTINUUM-RESUME-ENGINE-PRESENT` (the engine present under `constitution/submodules/continuum/` , `go test -race ./... green` , `continuum selfcheck` = good=PASS/bad=FAIL/negctrl=PASS, zero project literals) + paired §1.1 mutation (strip the literal → propagation gate FAILS; downgrade the resume path to read the ledger, or the oracle to pass its golden-bad → the mechanism gate FAILS; gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.207. Non-compliance is a release blocker. No escape hatch — no `--resume-reads-history` , `--skip-integrity-check` , `--unvalidated-oracle-OK` , `--hand-typed-resume-state` , `--per-project-resume-reimpl` flag.

§11.4.208 — Operator-request-history document: every project maintains a project-local, always-in-sync ledger of every operator request/prompt (User mandate, 2026-07-15). Verbatim operator mandate: introduce a request-history document capturing per operator request/prompt — full content, accepted-

timestamp with explicit timezone, the track that processed it, the alias that took the workable item, and the model + effort used — always in sync, no false data, no bluff; the operator's load-bearing correction (§11.4.35/§11.4.17): the RULE is universal (→ constitution submodule), the request-history DOCUMENT itself is PROJECT-LOCAL (never placed into the constitution). Every governed project MUST maintain a single project-local, append-only, newest-first **operator-request-history document** at a fixed project-declared path (e.g. `docs/requests/history.md` ; §11.4.35, never moved without a §11.4.66 operator decision), each entry carrying EXACTLY: (1) **Request content** (verbatim where archived, else a faithful complete summary, marked which); (2) **Accepted (when)** — date + time + TIMEZONE stated EXPLICITLY on every entry (consuming project declares its default zone §11.4.35; date-only ⇒ `time UNKNOWN` , never fabricated); (3) **Track** (§11.4.176/.178/.182); (4) **Alias** (§11.4.182/.196); (5) **Model + effort**. (A) **NO INVENTION (§11.4.6)** — an unrecoverable field is recorded literally `UNKNOWN` , never guessed; a fabricated timestamp/alias/model is a §11.4/§11.4.1 bluff at the requirements-record layer. (B) **HONEST RECONSTRUCTION BOUNDARY** — pre-mandate sessions are best-effort reconstructed from the durable dated record with unrecoverable fields marked `UNKNOWN` and the boundary stated explicitly, never silently implied complete. (C) **ALWAYS IN SYNC + EXPORTED** — §11.4.44 revision header, §12.10 append-on-every-new-prompt, §11.4.65/§11.4.73 four-format export (`.md + .html + .pdf + .docx` where applicable, siblings ≥ `.md`). (D) **KEEP-APPLYING MECHANISM** — a helper script and/or a `UserPromptSubmit` -class hook appends a newest-first row per NEW prompt at capture time (track/alias derived deterministically §11.4.182, honest `? / UNKNOWN`); a helper landed WITHOUT the auto-capture hook is honestly a partial mechanism (hook wiring = tracked §11.4.197/§12.10 follow-up, never claimed as auto-capture it does not perform). (E) **COMPLEMENTS, NEVER REPLACES** any per-session no-loss ledger (§11.4.197/.202), the workable-items DB SSoT (§11.4.93/.95), or the §12.10/§11.4.131 resumption artefacts; requirement loss at intake is FORBIDDEN (§11.4.197). (F) **RULE-UNIVERSAL / DOCUMENT-PROJECT-LOCAL (§11.4.35)** — the RULE lives in the constitution submodule; the document + its path + default timezone are consumer-owned project DATA (§11.4.28 decoupling), NEVER placed into the constitution. Classification: universal (§11.4.17). Composes §11.4.6/.35/.44/.65/.73/.118/.131/.176/.178/.182/.196/.197/.202/.93/.95/§12.10. Propagation gate `CM-COVENANT-114-208-PROPAGATION` (literal `11.4.208`) + recommended gate `CM-REQUEST-HISTORY-DOC-PRESENT` + paired §1.1 mutation.

Canonical authority: constitution submodule `Constitution.md` §11.4.208. Non-compliance is a release blocker. No escape hatch — no `--request-history-optional` , `--skip-request-log` , `--no-timezone-stamp` , `--invent-missing-field` , `--request-history-in-constitution` , `--prompt-without-history-entry` flag.

§11.4.209 — Code-review MUST run on the Fable model at xhigh effort (Opus xhigh fallback) (User mandate, 2026-07-15). Verbatim operator mandate: the mandatory independent code-review MUST ALWAYS be performed with the **Fable model at xhigh effort**; if Fable is unavailable, use **Opus at xhigh effort**. Every mandatory independent code-review the constitution requires — the §11.4.125 code-review-agent gate before pre-build + main build, the §11.4.142 universal every-change-review (no exception), the §11.4.194 exhaustive all-scenario/all-angle review, iterated to a zero-finding GO per §11.4.134 — MUST be dispatched on the **Fable model at xhigh effort**. (A) **Fallback (closed, ordered)** — ONLY when Fable is genuinely unavailable (established as FACT §11.4.6, never a guessed unavailability) the review runs on **Opus at xhigh effort**; a lower-effort or a non-Fable/non-Opus model is NOT an acceptable review substrate, `xhigh` is REQUIRED in both cases (the §11.4.194 all-angle depth is unreachable lower). (B) **Independence preserved (§11.4.70/.20)** — the Fable/Opus-xhigh model is the structurally-separated reviewer, never the author; §11.4.209 sets the review's MODEL + EFFORT, not its independence / iterate-until-GO (§11.4.134) / captured-evidence (§11.4.5/.69/.107). (C) **Honest boundary (§11.4.6)** — a review claimed done that ran on a weaker model / lower effort than Fable-xhigh (or Opus-xhigh on genuine Fable-unavailability) is a §11.4 bluff at the review-substrate layer; the review's captured evidence records which model + effort actually performed it (+ the Fable-unavailability FACT on fallback). STRENGTHENS §11.4.125/.142/.194/.134 — does NOT replace them; it pins the review's execution substrate. Classification: universal (§11.4.17). Composes §11.4.125/.134/.142/.194/.6/.20/.70/.196.

Propagation gate `CM-COVENANT-114-209-PROPAGATION` (literal `11.4.209`) + recommended gate `CM-CODE-REVIEW-FABLE-XHIGH` + paired §1.1 mutation.

Canonical authority: constitution submodule `Constitution.md` §11.4.209. Non-compliance is a release blocker. No escape hatch — no `--review-any-model` , `--review-lower-effort-OK` , `--skip-fable` , `--non-xhigh-review` , `--fallback-without-fable-unavailability-proof` flag.

§11.4.210 — Zero-loss request/prompt intake: every operator request/prompt MUST be mechanically CAPTURED, TRACKED, and PROCESSED — never skipped, ignored, avoided, or lost (User mandate, 2026-07-15). Verbatim intent: "every request/prompt we issue MUST ALWAYS be respected, taken into account, executed and processed! No avoiding, ignoring, skipping or any form of loss!" Every operator request/prompt MUST be (a) CAPTURED in the §11.4.208 request-history ledger MECHANICALLY at intake via a WIRED verified-runtime-state `UserPromptSubmit` auto-capture hook (NOT a manual step, NOT a "tracked follow-up" — §11.4.210 PROMOTES §11.4.208(D)'s hook from partial-follow-up to MANDATORY per the §11.4.205 enforced-not-advisory pattern; forensic FACT 2026-07-15: an entire release-drive session's prompts went un-recorded because the hook was left un-wired); (b) TRACKED → a §11.4.202 workable item / §11.4.197 completion, or an explicit evidence-backed no-op-with-reason; (c) PROCESSED — respected/executed, never dropped. Loss at ANY of the three stages is a §11.4 PASS-bluff at the requirements-intake layer of §11.4.197's severity class. Docs-chain-bound (§11.4.106), gated (§11.4.75 — hook verified-installed in the live hook path + content-matched, ledger fresh); the autonomous loop (§11.4.126/.87/.94/.97/.103) processes every captured request (done-condition never satisfied while any captured request is un-processed, §11.4.147(e)); honest §11.4.6 (UNKNOWN never invented). Classification: universal (§11.4.17). Composes §11.4.197/.202/.208/.205/.106/.126/.75/.6/.147(e). Propagation gate `CM-COVENANT-114-210-PROPAGATION` (literal `11.4.210`) + recommended gate `CM-REQUEST-CAPTURE-HOOK-WIRED` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.210. Non-compliance is a release blocker. No escape hatch — no `--skip-request-capture`, `--manual-ledger-OK`, `--hook-optional`, `--drop-unprocessed-request`, `--ledger-may-be-stale` flag.

§11.4.211 — Merge-conflict resolution during main → feature/product/flavor merges MUST run on the Fable model at xhigh effort (Opus xhigh fallback) (User mandate, 2026-07-16). Verbatim operator mandate: "During the merge of main branch to all feature/product/flavor branches of tracks 2 - 4, if conflicts happen they MUST BE resolved using Fable model with xhigh effort! If Fable is not available, then with Opus with xhigh!" Any merge-CONFLICT resolution performed while integrating the canonical trunk (`main` / `master`) into a feature / product / flavor branch — the §11.4.188 regular main → feature merge cadence, the §11.4.195 branch-taxonomy integration, the §11.4.41 merge-first pipeline, the

§11.4.167(D) trunk-into-stream sync AND the §11.4.167(I) approval-gated back-merge — MUST be carried out on the **Fable model at xhigh effort**; the merge-resolution analogue of §11.4.209 (which pins the same model + effort for code-REVIEW), pinning it for the DECISION-MAKING of conflict resolution (which side to keep, how to UNION both sides without loss, marker removal, semantic reconciliation of divergent edits). **(A) Fallback (closed, ordered):** ONLY when Fable is genuinely unavailable (not configured / rate-limited / unreachable, established as FACT §11.4.6 — never a guessed unavailability) the conflict resolution runs on **Opus at xhigh effort**; a lower-effort or a non-Fable/non-Opus model is NOT an acceptable substrate for resolving merge conflicts, **xhigh REQUIRED** in both cases (a careless resolution silently drops a commit or commits a marker). **(B) Scope:** the operator named tracks 2–4 (`feature / product / flavor`); GENERALISES per §11.4.17 to EVERY controlled main → (feature|product|flavor) merge on EVERY track that hits a conflict. **(C) Safety preserved:** sets the MODEL + EFFORT of the conflict-resolution work only, BOUNDED BY §11.4.113 (NEVER force-push) / §9 / §9.2 (no commit lost, hardlinked pre-op backup before any large/risky merge) / §11.4.41 step 3 (ZERO conflict markers committed, ZERO file silently dropped, union of both sides preserved) / §11.4.188(4) anti-bluff (post-merge smoke GREEN + marker-scan empty + no-lost-commit) / §11.4.84 (quiescent-only) / §11.4.37/§11.4.71 (fetch-first). **(D) Honest boundary (§11.4.6):** a merge whose conflicts were resolved on a weaker model / lower effort than Fable-xhigh (or Opus-xhigh on genuine Fable-unavailability) is a §11.4 bluff at the merge-resolution-substrate layer — the merge's captured evidence records which model + effort resolved the conflicts (and, on fallback, the FACT of Fable's unavailability). STRICT EXTENSION of §11.4.209 to merge-CONFLICT resolution; STRENGTHENS §11.4.188/§11.4.195/§11.4.41 by pinning the substrate of their conflict-resolution step. Classification: universal (§11.4.17). Composes §11.4.209/.188/.195/.41/.113/§9/§9.2/.6/.167(D)(I)/.84/.37/.71/.196. Propagation gate `CM-COVENANT-114-211-PROPAGATION` (literal `11.4.211`) + recommended gate `CM-MERGE-CONFLICT-FABLE-XHIGH` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.211. Non-compliance is a release blocker. No escape hatch — no `--merge-conflict-any-model`, `--skip-fable-merge-resolution`, `--non-xhigh-merge-conflict`, `--resolve-conflict-lower-effort-OK`, `--fallback-without-fable-unavailability-proof` flag.

§11.4.212 — Main README is the canonical starting-point / entry point for ALL project documentation — no doc may be an orphan unreachable from README (User mandate, 2026-07-16). Verbatim operator mandate: "Manual tests catalog must have linking all through the main README as the starting point of this and any other Project documentation! This fact and mandatory rule MUST be added into the constitution." The main README (`README.md` at the repository root) is the canonical STARTING POINT / entry point for ALL project documentation, for THIS project and EVERY project incorporating this constitution. Every project document in the §11.4.65 export scope — INCLUDING the manual-tests catalog (the §11.4.185 manual-QA testing catalog), the §11.4.153 per-feature Status / Status_Summary docs, the Issues / Fixed trackers + their summaries (§11.4.12 / §11.4.53 / §11.4.57), and every guide / research note / changelog / script companion doc — MUST be reachable / linked (DIRECTLY or TRANSITIVELY) from the main README. No project document may be an ORPHAN — a doc that no link-path from README reaches is a §11.4.212 violation. STRICT GENERALISATION of §11.4.57 (which mandated a `Tracked-Items + Status Documents` doc-link section covering the tracker / status set) to EVERY §11.4.65-scope document, so the README is the single navigable root of the project's documentation tree. **(A) Reachability = a link path.** README → guide → sub-doc is fine (transitive); the requirement is that SOME path from README reaches every in-scope doc, not that every doc is linked directly at top level. **(B) Always-sync (§11.4.59 / §11.4.106 docs-chain).** A newly-added §11.4.65-scope doc that no README-reachable path links is a §11.4.212 violation the moment it lands; the manual-tests catalog specifically MUST be linked from the README (the operator's load-bearing example). **(C) Honest boundary (§11.4.6).** Reachability is a link-graph property (every doc FINDABLE from README), NOT a claim about each doc's CORRECTNESS or freshness (those stay §11.4.44 / §11.4.106 / the doc's own gates). **(D) Composition with §11.4.57.** §11.4.57's `Tracked-Items` section remains the mandated README landing block for the tracker / status set; §11.4.212 adds that EVERY OTHER in-scope doc is likewise reachable — §11.4.57 is the specialisation, §11.4.212 the generalisation. Classification: universal (§11.4.17) — the consuming project supplies its README path, its manual-tests-catalog path, and its §11.4.65 doc-set enumeration per §11.4.35. Composes §11.4.57 / §11.4.59 / §11.4.65 / §11.4.106 / §11.4.153 / §11.4.185 / §11.4.44 / §11.4.12 / §11.4.53. Propagation gate `CM-COVENANT-114-212-PROPAGATION` (literal `11.4.212`) + recommended gate `CM-README-DOC-ENTRYPOINT-COMPLETE` + paired §1.1 mutation. **Canonical authority:**

constitution submodule Constitution.md §11.4.212. Non-compliance is a release blocker. No escape hatch — no `--orphan-doc-OK` , `--readme-not-entripoint` , `--skip-readme-link` , `--doc-unreachable-OK` flag.

§11.4.213 — FEATURE research-scheduling directive: recognized via all §11.4.140 forms, SCHEDULES (never synchronously executes) a deep, enterprise-grade research + implementation-planning effort as a tracked workable item (operator-directed mandate, 2026-07-16). Compact summary: extends the §11.4.202 reporting-directive family with a fourth registered action `FEATURE` — recognized via ALL SIX §11.4.140 grammar forms (`FEATURE :: x / DEFAULT::FEATURE :: x / /FEATURE x / /DEFAULT::FEATURE x / FEATURE ---> x /` the single-colon `FEATURE: x` , registered-action-only so lowercase prose like "feature: ..." never expands and is never asked about, §11.4.6) — that SCHEDULES rather than synchronously executes a deep research + implementation-planning effort: the directive creates a `Type=Task / Status=Queued` workable item (§11.4.15/.16/.54) by INVOKING the SAME §11.4.202 `report_item.sh` engine (never a duplicate implementation of the item-creation/DB-sync/tracker-push machinery — a second divergent tracker-push path would itself be a bluff-surface regression) and appends it to a durable `docs/requests/feature_queue.md` queue (mirroring `BACKGROUND` 's `docs/requests/background_queue.md` pattern) so a scheduled request is never dropped; the item's comprehensive description (§11.4.148/.171) embeds the FULL research mandate as its acceptance-defining work program — deep web research on best incorporation (§11.4.8/.99/.150), systematic-debug of every enumerated weak-spot/gap/danger-zone with a risk-free rock-solid design per gap (§11.4.102), always enterprise-grade/bleeding-edge/innovative, exhaustive documentation down to lines-of-code + micro-POCs + diagrams + SQL + templates (§11.4.65/.73), reuse-first investigation of vasic-digital/HelixDevelopment components kept decoupled (§11.4.28/.74/.177), full test-type + Challenges + HelixQA planning (§11.4.27/.169), fine-grained phased/task/subtask breakdown, enterprise-scalability + max-performance planning, OpenDesign UI/UX wireframes where applicable (§11.4.162/.190), full CodeGraph integration (§11.4.78/.79/.80), and fully-synced workable items across the SQLite SSoT + every connected external tracker (§11.4.93/.95/.148 D5), honestly SKIPPING any absent tracker (§11.4.10) — NEVER a faked push. The multi-day research itself is EXECUTED by the standing autonomous loop (§11.4.87/.94/.97/.103/.126) when it claims the item, driven to a genuinely COMPLETED-and-wired or explicitly evidence-backed CLOSED terminal

state under §11.4.197 — never left un-wired in the backlog; "scheduled" is never reported as "done" (§11.4.6). Anti-bluff four-layer coverage: pre-build gate `CM-FEATURE-DIRECTIVE` (FUNCTIONAL — sources + RUNS the §11.4.140 parser, never grep-only, §11.4.108) + paired §1.1 mutation + a real-DB end-to-end evidence trail (temp SQLite DB, no mocks, §11.4.27). Classification: universal (§11.4.17). Composes

§11.4.140/.202/.197/.8/.99/.150/.102/.65/.73/.28/.74/.177/.27/.169/.162/.190/.78/.79/.80/.93/.95/.14

Propagation gate `CM-COVENANT-114-213-PROPAGATION` (literal `11.4.213`) + recommended gate `CM-FEATURE-DIRECTIVE` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.213. Non-compliance is a release blocker. No escape hatch — no `--feature-runs-synchronously`, `--skip-feature-queue`, `--feature-without-research-mandate`, `--reimplement-tracker-push`, `--hardcode-project-in-engine`, `--skip-item-creation` flag.

EXTENSIONS landed 2026-07-17 (existing anchors strengthened — full text in `Constitution.md` ; research-derived from a consuming project's status-without-custody forensics). Forensic FACTs (genericised, all measured 2026-07-17): 576 tracked items, 182 claiming a done/ready status, **173 (95%) with NO registered regression guard** — an operator manually sampled 6 done-claiming items and found **6/6 broken**, the expected draw of that arithmetic, not bad luck; recorded reopen-per-recorded-fix **52%** (published elite <10%), itself an UNDERCOUNT because recurrences re-entered the tracker as NEW ids and statuses moved with ZERO history rows (one item: 4 Reopened events, 0 terminal events); the standing guard suite's exit code read ONLY its FAIL count — measured live on a release-candidate target: **61 registered guards → PASS 0 / FAIL 5 / PENDING 56 → exit 0 "not blocked"** — and the release-tag tooling consulted the guard suite **ZERO times**, so to the release gate "never verified" and "verified good" were the same colour; the project's own history selected the discriminator — every fix whose done-claim was preceded by an on-device RED → GREEN polarity flip that HELD recorded zero reopens, and every fix confirmed at source-green bounced. **No new anchor number was minted:** §11.4.115 and §11.4.146 already said what the operator mandates and did not bind because they are prose consulted by agents, not conditions checked by seams — the deliverable is the MECHANICAL BINDING (the §11.4.205 test applied: "if I documented this and nobody implemented the hook, would anything catch it?"). **§11.4.115(F) — machine-written verdict pairs + guard-viability:** RED and GREEN are harness-

written verdict files, never prose — carrying the item id, guard identity, polarity, exit code, the target's artifact fingerprint READ FROM THE TARGET AT RUN TIME, iterations (≥ 3 , §11.4.50), and evidence files whose CLASS matches the defect's layer (a source-grep transcript can never satisfy a user-visible/pixel/audio class — the wrong-layer oracle dies by construction); GREEN requires a prior RED with exit $\neq 0$ on the PRE-fix artifact and a DIFFERENT fingerprint (identical fingerprints prove the fix was never deployed); a guard never observed FAILing on the genuinely-broken artifact is unvalidated instrumentation and mints no verdicts (a guard structurally unable to FAIL — or to PASS — satisfies nothing); the canonical §1.1 mutation for a landed fix is the fix-commit's REVERT, and a mutation whose diff only deletes the literal strings the gate greps for is a refused tautology.

Recommended gate `CM-GUARD-VERDICT-MACHINE-WRITTEN` . **§11.4.135 (verdict-coverage extension) — ABSENCE of a verdict blocks exactly as a FAIL does:**

the suite's exit semantics MUST incorporate COVERAGE, never only the FAIL count; the release seam refuses the tag while `uncovered = registered ^ topology-present ^ no-verdict-for-the-candidate-fingerprint` is non-empty OR any FAIL exists; the release-tag tooling MUST consult the suite/verdict store mechanically; REQUIRES-REBUILD PENDING is honest on intermediate artifacts and self-clearing on the candidate (re-run on the candidate, or the item leaves the release's done-set — both forced, neither a permanent exemption); topology-absent guards are exempt ONLY via a checked-in target-pool topology map and MUST be enumerated in the release changelog as honest gaps (§11.4.3/ §11.4.118) — never silent; a naive "PENDING always blocks" is REJECTED as a §11.4.201 FAIL-bluff; pre-existing unguarded backlogs adopt via a ONE-TIME monotone-decrease ratchet snapshot so day one is green and the disease cannot spread. Recommended gate `CM-RELEASE-VERDICT-COVERAGE` (self-tested with golden-good / golden-bad / negative-control per §11.4.205/§11.4.107(10)).

§11.4.146(D3) — status-custody mechanical binding: a done/ready status write is REFUSED without the file chain `registry` row keyed by the item's EXACT stable id (checked-in alias map for legacy keys — a guard keyed to a non-item id fails the join and satisfies nothing) $\rightarrow$ `executable registered guard` $\rightarrow$ §11.4.115(F) RED+GREEN verdict pair $\rightarrow$ `class-matched evidence` — prose appears nowhere in the chain; Seam A = the workable-items engine (§11.4.93) refuses the write (project-agnostic per §11.4.28, custody paths are consumer DATA per §11.4.35); Seam B = a FULL-TABLE sweep gate on every pre-build run (not a diff check — catches raw-DB writes that bypass Seam A); Seam C = the §11.4.135 release seam; §11.4.205 construction order — the

executable sweep + its golden-bad meta-fixture (a temp SSoT overlay with a chain-less terminal item; the meta-test RUNS the sweep and requires FAIL, and asserts the sweep is WIRED) land in the SAME commit as any governance text citing them. Recommended gate `CM-STATUS-CUSTODY`. Honest boundary (§11.4.6), carried explicitly: the binding proves no status outruns its evidence and the REPORTED defect cannot silently return under the guard's conditions — it does NOT prove the feature bug-free (family dedup §11.4.186 remains), does NOT make a miscalibrated oracle honest (§11.4.107(10) goldens are necessary, not sufficient), does NOT replace the §11.4.185 manual-QA sufficiency gate, and makes target availability a HARD dependency of status progress BY DESIGN (a visible block replaces a silent pass). All Classification: universal (§11.4.17) — no hardware/vendor/project literal; consumers supply their registry, verdict-store, sweep, ratchet, and topology-map paths as DATA per §11.4.35. Propagation gates for the literals `11.4.115` / `11.4.135` / `11.4.146` are unchanged and stay GREEN; the three recommended mechanism gates + paired §1.1 mutations are gate-code = separate work item.

EXTENSIONS + NEW ANCHORS landed 2026-07-17 (round 2 — the carrier/ measurement-integrity harvest). Research-derived from one cycle's forensics in a consuming project, where EVERY false finding had ONE shape: **a token that MENTIONS X was matched as X, or an instrument returned silence that was read as data.** Genericised FACTs: a removal's completeness was measured by counting references to the removed component and the count matched the AUDIT-TRAIL COMMENTS documenting the removal (a completed removal read as failed); a pattern's escapes were eaten by an intervening shell so the far side searched a LITERAL and matched nothing (a false "the component is missing" that would have implied the target could not boot — it had booted); `\|` under an extended-regex dialect is a LITERAL pipe so a mandated-fields check searched a string nobody wrote and nearly reported the fields absent (all were present); a `check | head` pipeline returned the PAGER's exit status so a parse-check reported "clean ✓" three times without checking anything; a line-start-anchored pattern missed an indented, colour-escaped line and read a present verdict as no-verdict; a binary extractor split a multi-byte character and reported a present string as absent. §11.4.196(D)'s process-carrier and §12.12's self-match had ALREADY anchored this shape TWICE and it kept recurring — that recurrence is the argument for the general rule. The self-demonstrating instance: a mutation-residue scan run over the governance repo flagged the very anchor that DEFINES the mutation markers, because that anchor quotes them in its rule text (proven a

carrier, not residue, by an identical HEAD-vs-worktree count). **The fix, proven in-session: the CONTROL NEEDLE** — an artifact check returned 0 hits; instead of reporting absence it ran a string KNOWN present through the SAME path, got 1 hit, and only THEN reported the zero as real.

§11.4.201 (6)(7)(8) — scope widened to EVERY load-bearing measurement + carrier-vs-thing + control needle + metric-validation. (6) clauses (1)-(5) bind guards (decision seams that refuse/allow) with a refusal-shaped taxonomy; (6)-(8) bind **every measurement whose result is load-bearing** (grep/count/parse/exit-code/scan/metric/artifact-inspection) with a MEASUREMENT-shaped taxonomy — a **FALSE-NULL** (a zero read as absence when the instrument could not see) and a **FALSE-MATCH** (a hit on a carrier) are §11.4/§11.4.1 bluffs at the measurement layer; the null is more dangerous because a broken instrument and a clean artifact return the identical quiet zero. (7)(a) MATCH STRUCTURE, NOT SUBSTRING — distinguish the THING from a CARRIER that MENTIONS it. (7)(b) **A NULL IS NOT EVIDENCE until a CONTROL NEEDLE proves the instrument can see** — same instrument, same path, same artifact, a known-present needle MUST return non-null; needle absent ⇒ the instrument is blind and the zero says NOTHING (report the blindness, never the absence §11.4.6). **The needle MUST share the certified query's LOAD-BEARING FEATURES** (same dialect constructs / anchoring / quoting / encoding): a bare-literal needle certifies only the layers a bare literal crosses — the (7)(c) dialect FACT refutes the naive form directly, since a literal needle sails through the same tool+path and would 'certify' a zero from a query still searching a string nobody wrote, whereas a needle carrying the SAME alternation dies with it. Earned inference is therefore BOUNDED: needle found ⇒ the path can see the needle's QUERY CLASS ⇒ a zero from that same class is evidence; a zero from a query using features the needle never exercised needs a needle of its own class. Generalises §11.4.115(F)'s validate-the-detector from guards to every measurement, PER-QUERY — **golden fixtures cannot substitute** (§11.4.146(D3): goldens necessary-not-sufficient), an instrument can pass every fixture on the authoring host and be blind on the real path. (7)(c) **THE PATH IS PART OF THE INSTRUMENT** — the tool PLUS every layer the query crosses (wrapper-shell quoting, regex dialect, pipeline exit status, anchoring, encoding); each returns a CLEAN CONFIDENT WRONG answer without crashing, so §11.4.1's script-bug clause (which covers crashes-into-FAIL) never fires and §11.4.67's parse-check (parse-time only) cannot see it — the control needle is the ONLY mechanism that tests the path. (8) **A METRIC MUST BE VALIDATED**

AGAINST THE DEFINITION OF DONE — construct the correct end-state and evaluate the metric there; if the correct end-state does not move it to target, the metric is INVALID and MUST NOT gate work (a metric the correct outcome REFUTES will, if enforced, drive work AWAY from correctness); honest boundary §11.4.6 — this constitution already ships one such gating metric (§11.4.135's adoption ratchet, blocking at the RELEASE-TAG seam) whose validation against the definition of done is OWED, not claimed, and which every consuming project operating a ratchet MUST TRACK as a §11.4.197 item (stated as the obligation it is — this anchor asserts NO project has already filed it; asserting an unverified tracker state would be the §11.4.6 violation the clause forbids). **Interim precedence:** until that validation lands the §11.4.135(5) ratchet KEEPS GATING — clause (8) does NOT disarm it (an unvalidated ratchet blocking a release is bounded + visible + operator-resolvable; disarming the only adoption mechanism would silently re-open the unguarded-backlog hole §11.4.135 closes — the §11.4.101 reversible-safe choice); clause (8) binds every NEW metric from this anchor forward, and the ratchet's validation as owed work — never as a licence to switch it off. Gate contract landed in the canonical file: CM-GUARD-ASSERTS-REAL-CONDITION EXTENDED to the measurement layer ((i) a reported absence cites a class-matched control-needle result; (ii) the needle's class matches the certified query's load-bearing features) + NEW recommended sibling CM-METRIC-VALIDATED-AGAINST-DONE (every gating/ranking metric carries a recorded definition-of-done evaluation reaching target) + paired §1.1 mutations (bare-literal needle on a dialect-dependent query → (ii) FAILs; needle stripped → (i) FAILs; metric whose correct end-state moves AWAY from target → FAILs). Gate-code = separate work item — the contract is landed, the code is NOT claimed shipped (§11.4.6).

§11.4.112 (5) — impossibility verdicts are BOUNDED and MUST say where the edge is. FACT: a TRUE verdict ("relocating another application's protected video to a secondary display *while that application's UI stays in place, without its cooperation* is impossible") LEAKED to an ADJACENT VIABLE goal (whole-task placement) for six weeks, while the platform had shipped a public permission-free SDK call for the adjacent goal several major versions earlier and its own connected-display docs PRESCRIBED that call. The anchor was ASYMMETRIC: clause (3) hardens a verdict against reopening while nothing bounded its REACH. Now every structurally-impossible verdict MUST (a) STATE ITS EXACT SCOPE (the narrowest claim the evidence supports — every load-bearing qualifier is part of the verdict; a qualifier dropped is a verdict widened), (b) ENUMERATE

THE ADJACENT GOALS IT DOES NOT COVER (unexamined neighbour ⇒ UNKNOWN + tracked follow-up, never silently absorbed), (c) BE CITATION-FENCED (citing it beyond its fence, in planning/investigation/design/status, without its OWN §11.4.150 research + proof is a §11.4.6 violation + a verdict-layer bluff of the class §11.4.199 names). Honest boundary — outside the fence ≠ viable, it means UNDECIDED, earned by its own evidence never by inheritance in either direction. Why a clause not a narrative: §11.4.150 already documented this class in prose and REQUIRES the research pass before a verdict is EARNED — the leak still happened, because §11.4.150 fires where a verdict is MINTED and this leak happened where one was CITED as a premise, crossing no closure seam. Gate contract landed: CM-WONT-FIX-STRUCTURAL-IMPOSSIBILITY EXTENDED (scope statement + adjacent-goals-not-covered enumeration required; citing a verdict outside its stated scope without that goal's own §11.4.150 proof → FAIL) + paired §1.1 mutations. Gate-code = separate work item (§11.4.6).

§11.4.120 (seam-placement) — a gate whose precondition the gated work itself produces is on the WRONG SEAM. §11.4.120's trigger is a gate that FAILs (a gate that RAN); its sibling is a gate that CANNOT RUN in the imposed ordering, and the response is not reconciliation but re-seating. FACT: work was ordered "run checks E1-E4, then implement P1 only if they pass" — but the gate's own text defined E2 as running ON a P1-gated build, so E2 could not exist before P1; the source of truth gated ENABLING (building behind a default-off flag was permitted), and the gate had been imposed on AUTHORSHIP. (a) if gate G requires artifact A and A is produced by the work G gates, G belongs at a LATER seam (activation/enablement/release), never authorship — imposing it there is a deadlock, not rigor; (b) RE-READ the source of truth before re-seating (the gate's own text names its real seam and governs over any restatement — a restatement that moved it from enabling to writing is a §11.4.6 misstatement; obey the source, never weaken the gate); (c) the forbidden responses are §11.4.120's own (no fake-pass/delete/tautology, no build-then-retro-declare, and the work is NOT abandoned as "blocked" when the gate merely sits on the wrong seam). Honest boundary — an unrunnable gate is NOT automatically mis-seated (§11.4.102 first); it may be correctly seated with a genuinely-missing precondition = the §11.4.69 artifact_not_yet_built NOT-YET-RUNNABLE state. §11.4.110's clash detector cannot catch this class (its input is a DIFF; an ordering instruction is not a diff). The "report the circularity, don't improvise" half is ALREADY covered by §11.4.6/§11.4.66/§11.4.101/§11.4.201(4) — this clause claims only the seam-placement rule. Gate contract landed: CM-GATE-

RECONCILED-NOT-FAKE-PASSED EXTENDED (an imposed gate whose precondition the gated work itself produces → FAIL naming the circularity, re-seat at activation) + a golden-FALSE fixture that MUST include a correctly-seated gate merely awaiting its artifact (§11.4.69 `artifact_not_yet_built`) which the gate MUST NOT fire on, else it becomes the §11.4.201(1) false positive + paired §1.1 mutation. Gate-code = separate work item (§11.4.6).

§11.4.69 (closed reason-set) — `artifact_not_yet_built` added. The NOT-YET-RUNNABLE class (check CORRECT, topology PRESENT, precondition artifact absent — fix not built/deployed) had NO legal code, so a §11.4.135 REQUIRES-REBUILD PENDING had to misuse a false code (`hardware_not_present` / `feature_disabled_by_config` — both untrue, a §11.4.6 misstatement) or fail the helper at call time. Reporting it FAIL is a §11.4.1 FAIL-bluff (sends teams to fix healthy code); PASS is a §11.4 PASS-bluff. Its seam-dependence is §11.4.135's, NOT flat — honest+non-blocking on an intermediate artifact, BLOCKING on the release candidate; a flat "always blocks" is REJECTED per §11.4.135 as itself a §11.4.201 FAIL-bluff. **Audited, NO new anchor:** the operator-candidate "BLOCKED as a third verdict state" is ALREADY fully covered — §11.4.135's 2026-07-17 verdict-coverage extension states it more precisely (seam-dependent, and it explicitly rejects the flat framing), §11.4.3 forbids the PASS, §11.4.201(1) forbids the FAIL, §11.4.116/§11.4.45 already carry the verdict vocabulary; only the reason-code was missing.

§11.4.214 (NEW) — Recurrence-links-not-mints: a defect that returns MUST reopen its existing item, never enter as a new id. FACT: recurrences of already-"fixed" defects re-entered as BRAND-NEW ids (≥5 chains, one three deep), so the ORIGINAL's reopen counter never incremented — and the §11.4.55 `reopens_count`, the exact signal §11.4.132(d)/§11.4.189 use to rank the most-fragile work for the deepest scrutiny, **stays quiet precisely on the items that keep breaking**; the recorded 52% reopen rate was a KNOWN UNDERCOUNT for this reason; ~15 items mapped to ~9 real root causes, so "fix all 15" would have forked duplicate/conflicting fixes on one cause. Before minting a NEW id from ANY intake path, check for an existing item describing the SAME defect → RESOLVE the match THROUGH its `duplicate-of` links to the CANONICAL CHAIN HEAD (§11.4.90's closure vocabulary includes `duplicate-of`, and clause (5) MANUFACTURES such items — reopening one would resurrect a non-canonical copy beside its live canonical item, forking one defect into two and crediting the §11.4.55 signal to the WRONG item), then LINK it, and REOPEN (§11.4.34) **iff the**

resolved head is in a TERMINAL state; multi-match $\Rightarrow$ resolve every match to its head and act on the SINGLE survivor, >1 distinct head or a broken/circular link $\Rightarrow$ UNDECIDED (ask / mint-with-link, never guess §11.4.6). **Open-original case**: a matched item still OPEN gets LINK-ONLY — 'reopen an open item' is undefined in the §11.4.15 status machine and §11.4.34's Reopened presupposes a closure to demote FROM (§11.4.7); the reopen path exists to correct the specific lie 'closed as fixed and it is not', and an open item is not telling it. **§11.4.202 precedence (explicit, never left to a consumer)**: §11.4.202's mandate is that no report is ever LOST — SATISFIED IN FULL by a SAME-DEFECT report landing as link + (terminal-only) reopen; it forbids the report evaporating into prose, it does NOT require a fresh id per report. §11.4.214 refines WHICH item a report lands on; §11.4.202 owns THAT it lands; DISTINCT/UNDECIDED verdicts run §11.4.202's mint path unchanged. (1) EVERY intake path binds (reporting directives §11.4.202, AI-detected test/gate failures, regression-guard verdicts, cycle re-discovery, operator manual testing — §11.4.34's own reason vocabulary enumerates them; wiring dedup ONLY into the reporting channel under-covers exactly the machine-driven paths that recur most). (2) Key on ticket else normalised (subject, scope) — NEVER a bare subject substring (§11.4.186). (3) **A wrong merge silently DELETES a real defect — so distinct-but-similar is a FIRST-CLASS verdict**: SAME-DEFECT $\Rightarrow$ link+reopen, DISTINCT $\Rightarrow$ mint (correct + expected), UNDECIDED $\Rightarrow$ ask (§11.4.66/.105) or mint-WITH-a-candidate-duplicate-link; autonomous default (§11.4.101) is mint-with-link over merge — a spurious id is cheap + recoverable, a wrongly-merged defect is LOST (the asymmetry decides). (4) Self-validated oracle (§11.4.107(10)/§11.4.186): golden-good + golden-bad + a **negative-control** (two genuinely DISTINCT same-subject items that MUST NOT merge) — an oracle without it is a false-merge engine. (5) Id stability preserved (§11.4.54 — reopen the original, never renumber/reuse/retire; an already-minted duplicate is LINKED via §11.4.90 duplicate-of, never silently deleted). (6) The restored signal IS the point — an undercounting reopen counter is a §11.4/§11.4.1 bluff at the PRIORITISATION layer: it reports the fragile set as stable and aims the deepest scrutiny at the wrong items. Honest boundary §11.4.6 — dedup makes the signal TRUE, not the item correct, and a same-symptom recurrence with a different root cause is real (§11.4.102 decides; the linked+reopened item is where that is recorded, never a reason to fork). Gate CM-RECURRENCE-LINKS-NOT-MINTS (every CONFIRMED-duplicate link whose target was TERMINAL at link time has a reopen event on that target) — **its own scope is §11.4.201-bound**: it MUST NOT fire on an OPEN-target link (LINK-only is correct

there) nor on a clause-(3) UNDECIDED candidate-duplicate link (un-reopened BY DESIGN, and the mandated autonomous default), or it becomes exactly the §11.4.201(1) false-positive refusal this same round harvests against; its golden-FALSE set MUST include both + paired §1.1 mutation. Why NEW: §11.4.55 CONSUMES the count, §11.4.34 presupposes the status is already Reopened , §11.4.54 forbids only one-id-two-items, §11.4.186's DEDUP needs DIVERGENT metadata and gates export/sync/commit not INTAKE, and §11.4.202's operative text, read WITHOUT this anchor, reads as mint-every-time — the reading the precedence paragraph above corrects (they are NOT in conflict: §11.4.202 owns THAT a report lands, §11.4.214 refines WHICH item it lands on).

§11.4.215 (NEW) — A doc that BINDS work MUST live, tracked, in the repository where that work happens. FACT: the document DEFINING a binding acceptance gate was UNTRACKED and lived in an out-of-tree scratch directory belonging to a different checkout — so the gate could not be honoured (unreadable), reviewed (no diff/history/authorship), version-controlled (no revision/rollback), or reconciled when it contradicted the source of truth. A binding artifact that lives nowhere is a rumour with an imperative mood. (1) **BINDING ⇒ TRACKED, and un-tracked ⇒ NOT BINDING** — a binding claim sourced from an untracked/out-of-tree doc MUST NOT be enforced; land the doc or decline the gate, never obey a citation of it. (2) **IN THE REPOSITORY THE WORK HAPPENS IN** — a sibling checkout / scratch dir / chat message / issue comment is NOT tracked for this purpose (the test: can someone who clones ONLY that repo read the binding text at its declared path?); multi-repo binding docs are inherited BY REFERENCE from a tracked shared dependency (§11.4.28/§11.4.177), never copied, never out-of-tree. (3) **The prior question §11.4.212 cannot ask** — §11.4.212 ENUMERATES the in-repo scope and checks reachability, so a doc in no checkout is never enumerated and its gate cannot FAIL; §11.4.212 answers "is it findable?", §11.4.215 answers "is it here at all?" (once in-repo, §11.4.212/.44/.65/.73/.106 all attach as normal — this is their precondition, not their competitor). (4) **The generalisation of §11.4.95** (whose entire content is "authoritative ⇒ tracked, NOT a build artefact, IS authoritative source data" for ONE named artifact), exactly as §11.4.212 was minted as the generalisation of §11.4.57. §11.4.11 is the counterweight and stays intact — logs/forensics/scratch remain out-of-tree and untracked by DEFAULT; **binding is what pulls a document into the repository**, and a doc nobody cites as binding is governed by §11.4.11, not this anchor. Honest boundary §11.4.6 — in-repo presence proves AVAILABLE + ATTRIBUTABLE +

VERSION-CONTROLLED, never CORRECT/current (§11.4.44/.106/.186), and never well-seated (§11.4.120's seam clause — whose forensic case was carried by this SAME out-of-tree document: two defects, one artifact, which is itself the argument that an unreviewable binding doc is where bad gates hide). Gate `CM-BINDING-DOC-IN-REPO` + paired §1.1 mutation.

All Classification: universal (§11.4.17) — no hardware/vendor/project literal; consumers supply their intake wiring, id scheme, normalisation function, and binding-doc paths as DATA per §11.4.35. Propagation gates `CM-COVENANT-114-214-PROPAGATION` + `CM-COVENANT-114-215-PROPAGATION` (literals `11.4.214` / `11.4.215`); the literals `11.4.201` / `11.4.112` / `11.4.120` / `11.4.69` are unchanged and their propagation gates stay GREEN. All recommended mechanism gates + paired §1.1 mutations = gate-code, a separate work item.

DESIGN-METHODOLOGY MANDATES landed 2026-07-22 (§11.4.216–§11.4.223 — full text in `Constitution.md` ; proving project: the Helix Thready MVP design package at `docs/public/research/mvp/design/`). The complete design methodology built and verified end-to-end in Helix Thready is codified as eight universal anchors so every governed project inherits it out of the box. All Classification: universal (§11.4.17) — consumers supply brand values, platform lists, paths, and submodule pointers as DATA per §11.4.35. Propagation gates `CM-COVENANT-114-216-PROPAGATION` ... `CM-COVENANT-114-223-PROPAGATION` (literals `11.4.216` – `11.4.223`) + recommended mechanism gates + paired §1.1 mutations = gate-code, a separate work item.

§11.4.216 — Canonical machine-readable design-token source (User mandate, 2026-07-22). Compact summary: every UI-shipping project maintains its design tokens as ONE canonical machine-readable source — CSS custom properties, light in `:root {}` , dark via the `[data-theme="dark"]` override (composing the sanctioned dark mechanisms) — and EVERY surface binds via `var(--...)` or a GENERATED binding produced by codegen FROM that file, never a hand-kept copy (hand-copies drift silently); NO INVENTED VALUES (§11.4.6 applied to design): every color/size/duration traceable to captured evidence — brand colors to a recorded eyedrop/measurement of the canonical brand asset (proving project: brand `#B6E376` = mean of 1,625,855 logo pixels), accent pairs to MEASURED WCAG ratios against their PAIRED background (6.03:1 light / 13.56:1 dark), structural values to the documented shared scale; a hex resolving to no

token is a violation — use the nearest token, never mint one; the token source is the white-label seam (§11.4.217 closed override set); not-yet-built codegen bridges are [OPEN]-tracked (§11.4.223), never claimed shipped. Composes §11.4.6/.162/.170/.190/.197/.217/.223. Propagation gate `CM-COVENANT-114-216-PROPAGATION` + recommended gate `CM-DESIGN-TOKEN-SINGLE-SOURCE` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.216. Non-compliance is a release blocker. No escape hatch — no `--hand-copied-palette-OK`, `--inline-hex-OK`, `--invent-token-value`, `--light-only-tokens`, `--platform-keeps-own-palette` flag.

§11.4.217 — OpenDesign brand contract: 9-section DESIGN.md + tokens.css twin, disjoint color roles, AA-pinned accents, locked attribution footer (User mandate, 2026-07-22). Compact summary: every UI-shipping product carries under `design/opendesign/` (or its declared equivalent §11.4.35) a `DESIGN.md` in the VERIFIED OpenDesign 9-section schema (H1 + > Category: blockquote immediately after the H1; number-prefix-parsed sections; a "Font labels for catalog extraction" block; real hex; CSS variables wrapped in `:root {}` with dark via `[data-theme="dark"]`; real component CSS with `prefers-reduced-motion` targeting specific elements, never a global `*`; `:focus-visible` on every interactive component — format verified against the tool's CURRENT source per §11.4.99, never from memory) + the machine-readable `tokens.css` twin (the §11.4.216 canonical file or a compiled view; prose and tokens never diverge); THREE DISJOINT color roles by contract — brand DECORATIVE only (never body text), accent INTERACTIVE + AA-PINNED (measured $\geq 4.5:1$ per mode), semantic STATE never overridden by any brand (a brand color can never mask an error signal); white-label = a CLOSED override set (brand tokens, accent pair, logo, slogan) with accent overrides SERVER-VALIDATED to WCAG AA or REJECTED carrying the measured ratio + an AA-passing suggestion; structural/neutral/semantic tokens NOT account-overrideable; the organization attribution footer is LOCKED — present on every surface, never removed even fully white-labeled. Composes §11.4.6/.99/.162/.190/.216/.220/.223. Propagation gate `CM-COVENANT-114-217-PROPAGATION` + recommended gate `CM-OPENDESIGN-BRAND-CONTRACT` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.217. Non-compliance is a release blocker. No escape hatch — no `--skip-brand-contract`, `--prose-only-design-md`, `--brand-as-text-OK`, `--unvalidated-accent-override`, `--remove-attribution-footer`, `--semantic-tokens-brandable` flag.

§11.4.218 — Living design-library catalogue + per-cell-provenance reusability matrix (User mandate, 2026-07-22). Compact summary: every UI-shipping project maintains a LIVING component catalogue rendering EVERY component × declared state × light/dark from SELF-CONTAINED artifacts (zero external requests — no CDNs, vendored libs inline, honest font fallbacks), token-bound (§11.4.216), with REAL interactive behaviors (real dialog/sort/toggle; `prefers-reduced-motion` honored per §11.4.221) — proving project: 14 groups, ~38 components, ~118 state cells, ~236 rendered states, both themes auditable via a real persisted toggle; PLUS a components × platforms reusability matrix where EVERY cell carries closed-set provenance — VERIFIED (re-checked at source, dated) / PROPOSED ([DEFAULT — adjustable]) / ASSUMED (labeled inference); an unmarked cell reads ASSUMED; a scaffold cell claiming VERIFIED is a §11.4 PASS-bluff at the design-planning layer; the matrix is read BEFORE implementing any component on any platform. Spec'd-but-unrendered components/states are tracked gaps (§11.4.223). Composes §11.4.6/.162/.170/.216/.221/.223. Propagation gate `CM-COVENANT-114-218-PROPAGATION` + recommended gate `CM-DESIGN-LIBRARY-LIVING-CATALOGUE` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.218. Non-compliance is a release blocker. No escape hatch — no `--cdn-in-catalogue`, `--single-theme-catalogue`, `--mocked-behaviors-OK`, `--unmarked-matrix-cell`, `--scaffold-claims-verified` flag.

§11.4.219 — Per-app screen catalogues: full-IA coverage, per-platform chrome, honesty markers (User mandate, 2026-07-22). Compact summary: every shipped surface (web / mobile / desktop / TUI / marketing / ...) has a screen catalogue covering EVERY node of its information architecture (an IA node with no catalogue screen is a tracked gap, never silently skipped); per-platform CHROME VARIANTS rendered where one design serves several hosts (mobile OS chrome toggle, per-OS window chrome, terminal-width constraints); each screen SELF-CONTAINED, token-bound (§11.4.216), light + dark, populated with REALISTIC product data and a per-screen state contract (loading/empty/error/populated); UNBUILT or UNDECIDED behavior is HONESTLY marked — `[GAP: id]` / `[OPEN: id]` (§11.4.223), labelled placeholders, inline notes, disabled controls — NEVER a designed-as-if-real fake flow (a fake flow presented as settled is a §11.4 PASS-bluff at the product-design layer). Proving project: 13 web + 11 mobile (Android/iOS chrome toggle) + desktop shell + 10 TUI + 3 marketing screens, every unbacked flow marked. Composes §11.4.6/.162/.170/.185/.190/.216/.218/.223.

Propagation gate `CM-COVENANT-114-219-PROPAGATION` + recommended gate `CM-SCREEN-CATALOGUE-FULL-IA` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.219. Non-compliance is a release blocker. No escape hatch — no `--partial-ia-coverage`, `--one-chrome-fits-all`, `--fake-flow-OK`, `--lorem-ipsum-content`, `--single-theme-screens` flag.

§11.4.220 — Open-first design tooling: self-hosted PenPot primary via the PenPot submodule; proprietary tools import/export only; in-repo open-format sources (User mandate, 2026-07-22).

Compact summary: the collaborative design platform is an open-source self-hostable tool — PenPot as the standing default — run SELF-HOSTED and consumed ONLY via the dedicated reusable PenPot submodule (working name `vasic-digital/PenPot`, final repository name pending — the generic referent is normative meanwhile) carrying ALL PenPot mechanisms/utilities (deployment recipe, RPC client, access-token minting, import tooling, browser-hydration automation, MCP wiring), itself consuming the Containers submodule (§11.4.76, rootless §11.4.161) and DECLARING that dependency in its §11.4.31 `helix-deps.yaml`; projects MUST NOT hand-roll PenPot deployment/client code (§11.4.28/§11.4.74 — extend the submodule upstream); the submodule carries mechanisms ONLY, never project design content. Proprietary tools (Figma or equivalent) are OPTIONAL import/export/review bridges, NEVER the source of truth — no design value may exist only inside a proprietary tool/cloud; design sources live IN-REPO in open, diffable formats (SVG masters, CSS tokens, JSON manifests/IR/variables payloads, Lottie JSON, self-contained HTML — §11.4.215 applied to design). Tooling capability claims are RUNTIME-PROBED (§11.4.201/§11.4.99) with a headless-vs-operator-interactive matrix + an honest-limits section — a cloud push claimed without a real credentialed run is a fabrication. Operational rules (research-verified): the HEADLESS write path is the token-authenticated RPC API — access tokens (minted under the `enable-access-tokens` backend flag, handled per §11.4.10) are the MANDATED automation auth mechanism, never interactive OAuth relays; API-created PenPot files REQUIRE a browser hydration pass before handoff (thumbnails/text layout materialize in the browser runtime, not the write API) — an unhydrated API-created file handed off as render-ready is a §11.4 PASS-bluff at the design-handoff layer. Composes

§11.4.6/.10/.28/.31/.69/.74/.76/.99/.161/.162/.201/.215/.216/.223. Propagation gate `CM-COVENANT-114-220-PROPAGATION` + recommended gate `CM-OPEN-DESIGN-`

TOOLING-PRIMARY + paired §1.1 mutation. **Canonical authority:** constitution submodule Constitution.md §11.4.220. Non-compliance is a release blocker. No escape hatch — no --figma-is-source-of-truth , --cloud-only-design , --proprietary-format-source , --hand-rolled-penpot-stack , --oauth-relay-automation , --skip-hydration-pass , --unprobed-tooling-claim , --rootful-design-stack flag.

§11.4.221 — Motion discipline: token-bound timing + machine-readable manifest + reduced-motion static fallbacks (User mandate, 2026-07-22).

Compact summary: every duration/easing binds to the §11.4.216 motion tokens (an ad-hoc timing literal bypassing the token layer violates §11.4.6); every playable animation ships a STATIC reduced-motion fallback DERIVED FROM ITS OWN KEYFRAMES (poster frame for loopers, final frame for one-shots) in an open format, named per animation in a versioned MACHINE-READABLE motion manifest (id, loop mode, poster, reduced-motion mode, reducedMotionFallback) — a deliberate instant-cut is an EXPLICIT null-with-reason, never a silent absence; prefers-reduced-motion: reduce honored on EVERY surface via targeted component gates (never a global *), swapping in the manifest-named fallback; NO animation is claimed working without a REAL render in a real runtime (parse + structural schema + rendered geometry with progression assertions — the §11.4.107(10) analyzer discipline applied to motion assets), unexercised runtimes [OPEN]-tracked. Proving project: 6 hand-authored Lottie animations + 5 keyframe-derived fallback SVGs + manifest v2, lottie-web/jsdom + headless-Chromium validated, mobile runtimes honestly [OPEN]. Composes §11.4.6/.107(10)/.162/.170/.216/.218/.223. Propagation gate CM-COVENANT-114-221-PROPAGATION + recommended gate

CM-MOTION-TOKEN-BOUND-REDUCED-FALLBACK + paired §1.1 mutation. **Canonical authority:** constitution submodule Constitution.md §11.4.221. Non-compliance is a release blocker. No escape hatch — no --ad-hoc-duration , --no-reduced-motion-fallback , --global-star-motion-kill , --manifest-optional , --unrendered-animation-claimed-working flag.

§11.4.222 — Design-completion export wave: a design is not done until the exports land, honestly (User mandate, 2026-07-22).

Compact summary: a design (system / brand / screen set / change wave) is NOT done until its EXPORT WAVE runs with verified outputs — (1) per-platform token codegen regenerated from the canonical source (or [OPEN]-tracked); (2) the FULL icon/asset raster matrix, documented per OS surface and locked 1:1 to its generator, produced by an

HONEST-OUTPUT pipeline: every candidate renderer PROBE-TESTED with a content-verified probe render before use (§11.4.201 — presence ≠ correctness), fall-through on probe failure, absent tools SKIP-with-note (§11.4.3/§11.4.69), NEVER a placeholder/blank/faked output, every raster verified non-degenerate (size floor + non-blank; multi-resolution bundles opened + verified per §11.4.38) — proving project: 83 real rasters, zero faked; (3) both-theme screen renders (deterministic theme forcing, every render verified non-blank) + the bound design book; (4) hand-off bundles with per-leg HONEST statuses (produced / deferred-operator-only / impossible-with-reason); (5) the §11.4.65/§11.4.73 doc twins regenerated. A design wave reported complete without its export wave — or with any faked/unverified export output — is a §11.4 PASS-bluff at the design-delivery layer. Composes §11.4.3/.6/.38/.65/.69/.73/.201/.216/.219/.220. Propagation gate `CM-COVENANT-114-222-PROPAGATION` + recommended gate `CM-DESIGN-EXPORT-WAVE-COMPLETE` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.222. Non-compliance is a release blocker. No escape hatch — no `--skip-export-wave`, `--unprobed-renderer`, `--placeholder-raster-OK`, `--single-theme-renders`, `--fake-bundle-status`, `--stale-doc-twins` flag.

§11.4.223 — Provenance-marker discipline for design docs + central open-item registry (User mandate, 2026-07-22). Compact summary: every design document carries inline provenance markers from a closed, project-declared vocabulary whose MINIMUM CORE is `[VERIFIED]` (with evidence/check date) / `[OPERATOR]` / `[DEFAULT – adjustable]` / `[OPEN: id]` (extensions: `[RESEARCH]`, `[ASSUMED]`, `[BUILD-NEW]`, `[GAP: id]`, `[IN-HOUSE: ...]`, `[CONSTITUTION §x]` per §11.4.35); an unmarked load-bearing claim reads UNVERIFIED, and an assumption presented as verified is a §11.4.6 violation of §11.4 PASS-bluff severity at the design-documentation layer; every unresolved decision is an `[OPEN: id]` with a stable, never-renumbered id (§11.4.54 discipline) owned by a named file; the design area index maintains the SINGLE CANONICAL open-item registry, TREE-VERIFIED in both directions (every tree id in the registry, no registry ghosts — the §11.4.186 consistency discipline applied to design opens), closures citing the closing revision, narrowed items recording exactly what closed and what remains; registry items map to tracked workable items (§11.4.15/§11.4.54/§11.4.93) so the §11.4.197 completion mandate applies — no design open is silently abandoned. Proving project: 41 registered ids across nine families, tree-verified, 2 closed with citations. Composes

§11.4.6/.15/.44/.54/.61/.93/.123/.186/.197/.212. Propagation gate CM-COVENANT-114-223-PROPAGATION + recommended gate CM-DESIGN-DOC-PROVENANCE-MARKERS + paired §1.1 mutation. **Canonical authority:** constitution submodule Constitution.md §11.4.223. Non-compliance is a release blocker. No escape hatch — no --unmarked-claims-OK , --assumption-as-verified , --unregistered-open-item , --registry-optional , --silently-abandon-open flag.

§11.4.224 — Test-first (TDD) for ALL work, not only fixes + a minimum code-coverage floor that is NECESSARY-never-sufficient (User mandate, 2026-07-22). Verbatim operator mandate: "Any work we do MUST START by writing the test! We MUST FULLY comply and follow the TDD - The test driven development! It MUST BE applied for any change, implementation, wiring, tests, challenges and any other form of codebase we write! Bash scripts as well! Coverage with TDD MUST BE close to 100%, not less than 85% !!!" **The gap closed:** §11.4.43's operative text binds FIXES only ("Every fix MUST follow the 5-step TDD-fix workflow"), so a NEW feature / wiring / gate / Challenge / BASH SCRIPT inherited no test-first obligation; and §11.4.27/§11.4.169's "100%" is test-TYPE breadth (which KINDS of test exist per feature × platform), stating no CODE-coverage figure anywhere — §11.4.224 closes exactly those two holes and claims nothing else as new. **(A) TEST-FIRST BINDS ALL WORK** — every work product with an executable surface (change, feature, implementation, wiring, refactor, gate, §1.1 mutation, Challenge, test-harness code, AND bash/shell scripts): the test is written FIRST, run FIRST, and OBSERVED TO FAIL for the right reason before the implementation exists (RED on the broken artifact per §11.4.115 where one exists; RED against the absent behaviour where the work is new); a test authored AFTER its implementation proves only that it AGREES with the code, never that it CATCHES the defect (the §11.4.43 PASS-bluff generalised). Per-work-class test named explicitly (§11.4.6): bash/shell script → an executing test through its REAL invocation path asserting exit status + output + state delta, NEVER a bash -n parse-check alone (parse-cleanliness is §11.4.67's separate obligation and proves nothing about behaviour); a pre-build GATE → its paired §1.1 mutation IS the test-first artifact, authored and observed to make the gate FAIL BEFORE the gate is trusted (§11.4.115(F) — observation-before-trust, NOT authorship order); a Challenge → its scored assertion against captured evidence (§11.4.5/§11.4.69); a behaviour-bearing config/data change (any committed data an executing consumer reads to decide behaviour) → an executing test that drives that CONSUMER

through the changed value and asserts the changed observable behaviour. **(B) CODE-COVERAGE FLOOR ≥ 85%, ~100% TARGET**, applying to the executable codebase INCLUDING bash/shell (the operator's "Bash scripts as well!"). **(C) NECESSARY, NEVER SUFFICIENT — anti-gaming (§11.4.201(8)/§11.4/§1.1):** MEASURED FACT (in-session, 2026-07-22) — a line-granular instrument reported an `if/else` line COVERED while the `else` branch was NEVER taken, and an assertion-free test raises the number identically to a proving one; therefore the REAL bar remains catch-its-own-negation proven by a paired mutation (a line executed by an assertion-free test is NOT covered in the sense that matters), the floor is a MINIMUM ON A PROXY, and citing "≥ 85% covered" as proof a change works is a §11.4 PASS-bluff at the metric layer; the floor NEVER stands alone (composes §1.1 + §11.4.5/§11.4.69/§11.4.107 captured evidence + §11.4.108 runtime signature). **(D) METRIC VALIDATION (§11.4.201(8)) — SPLIT VERDICT, honestly stated:** directional validation PERFORMED 2026-07-22 (at the correct end-state — every behaviour test-driven, every test catching its negation — the implementing lines ARE executed by their tests, so the metric DOES reach target; the floor is VALID as a NECESSARY condition); the CONVERSE is REFUTED by the clause-(C) measured gameability fact; and the EMPIRICAL per-corpus calibration of the 85% threshold is explicitly OWED, NOT claimed — every consuming project enforcing the floor MUST track it as a §11.4.197 item (this anchor asserts no project has already filed it — asserting an unverified tracker state would be the §11.4.6 violation the clause prevents); INTERIM PRECEDENCE mirroring §11.4.201(8)'s own §11.4.135-ratchet precedent — the floor KEEPS GATING while the calibration is owed (an uncalibrated floor blocking a merge is bounded, visible, operator-resolvable; disarming it silently re-opens the untested-code hole — the §11.4.101 reversible-safe choice), never a licence to switch it off. **(E) HOW MEASURED, per language, honestly (§11.4.6/§11.4.35):** the consuming project supplies its per-language instrument as DATA; the universal rule is that an instrument MUST exist and MUST actually be RUN — a floor nobody can measure is an unmeasurable gate and itself a §11.4.201 bluff, and a percentage nobody measured is a §11.4.6 fabrication. MEASURED host reality (2026-07-22): Go built-in coverage (`go test -coverprofile + go tool cover`) + `gcov` PRESENT, but `kcov` / `bashcov` / `bats` / `shellspec` / `shunit2` ALL PROBED ABSENT for bash (`shellcheck` IS present but is a LINTER, not a coverage instrument, and satisfies nothing here). PROVEN-FEASIBLE zero-tooling bash mechanism, measured in-session never asserted: `PS4='+COV:${BASH_SOURCE##*/}:${LINENO}:' bash -x <script>` emits executed line numbers, and executed ÷

executable (non-blank, non-comment) source lines yields line coverage — a probe correctly identified an uncovered function body, and subshell bodies ARE traced. HONEST LIMITS (measured, stated as fact): LINE coverage NOT branch coverage (the clause-(C) same-line `if/else` counted covered with one branch never taken); `set +x` regions and traps unaccounted — a FLOOR-measuring instrument only, so branch-level assurance requires a real tool (e.g. `kcov`) or an honestly-recorded shortfall. Where NO instrument exists for a language: INSTALL one, or record an §11.4.3/§11.4.69 SKIP-with-reason + a tracked §11.4.197 item — NEVER an invented percentage, NEVER a silently-dropped floor. EXCLUSION-LIST FENCE (amendment, remediation round of 2026-07-22 — the §11.4.135 exemption pattern applied to the coverage corpus): a corpus exclusion is legal ONLY via a CHECKED-IN exclusion list, every entry justified from the closed class set {generated-code (output of a committed generator per §11.4.77) | vendored-third-party | non-shipping-fixtures-and-golden-assets} and enumerated as honest gaps never silent; excluding FIRST-PARTY executable code additionally REQUIRES a tracked §11.4.197 item naming the plan to bring it into scope — an unlisted, unjustified, or un-tracked-first-party exclusion VOIDS the measured figure (a consumer that excludes most of its first-party executable code and truthfully measures $\geq 85\%$ on the remainder has satisfied the floor in letter and voided it in substance — the silent-evasion channel this fence closes). **(F) HONEST**

BOUNDARY / SCOPE (§11.4.6): "all work" = work products with an EXECUTABLE surface; a pure prose/governance/documentation edit has no executable surface and is NOT coverage-scoped (its discipline is §11.4.44/§11.4.65/§11.4.106/§11.4.212 + the §11.4.142 universal review — demanding a coverage percentage there would be a §11.4.201(1) false-positive refusal), BUT a governance change that SHIPS an executable artifact (gate, hook, script, mutation) is FULLY in scope for (A) and (B); test-first PRECEDES §11.4.108 four-layer verification / §11.4.40 full-suite retest / §11.4.185 manual-QA, never replaces them; and the CODE-coverage axis is DISTINCT from §11.4.27/§11.4.169's TEST-TYPE breadth — both apply, neither substitutes, satisfying one is never evidence for the other. Classification: universal (§11.4.17) — the consuming project supplies its per-language instruments, corpus scope, CHECKED-IN clause-(E)-fenced exclusion list, and threshold calibration per §11.4.35. Composes

§11.4/.1/.5/.6/.27/.35/.43/.50/.66/.67/.69/.77/.85/.107/.108/.110/.115/.122/.135/.142/.146/.169/.185/.

§1.1. Propagation gate `CM-COVENANT-114-224-PROPAGATION` (literal `11.4.224`)

+ recommended gates `CM-TEST-FIRST-ALL-WORK` + `CM-COVERAGE-FLOOR`

(contract amended, remediation round of 2026-07-22: the measured numerator

counts ONLY RED-CAPABLE tests — every contributing test carries a recorded failing-first run OR a paired mutation observed to make it FAIL, zero-assertion contributors rejected by construction, sampling-based mutation verification an acceptable stated implementation — so ONE real RED test plus assertion-free padding to 85% FAILs the gate; and the gate FAILs on any exclusion that is unlisted, outside the clause-(E) closed class set, or first-party without its tracked §11.4.197 item; brownfield adoption is DELIBERATELY UNDEFINED — the owed gate-code item MUST surface the adoption question to the OPERATOR per §11.4.66 with concrete options {immediate hard floor | one-time §11.4.135-style monotone-decrease ratchet | per-corpus phase-in | changed-code-only with a scheduled full-corpus deadline} BEFORE first enforcement, recording the answer as consumer DATA — never an invented ratchet, a §11.4.122-class decision the operator owns) + paired §1.1 mutations; gate-code = a separate work item — the CONTRACT is landed, the gate CODE is NOT claimed shipped (§11.4.6).

Canonical authority: constitution submodule `Constitution.md` §11.4.224. Non-compliance is a release blocker. No escape hatch — no `--test-after-OK`, `--skip-red-first`, `--bash-exempt-from-tests`, `--coverage-not-applicable`, `--below-floor-OK`, `--coverage-proves-correct`, `--unmeasured-percentage-OK`, `--assertion-free-test-counts` flag.

§11.4.225 — Scheduler-quota burst-throttling telemetry + interactive-scope isolation from bursty fleets (research-derived, 2026-07-23). Compact summary: progressive / bursty / load-correlated interactive degradation MUST be diagnosed from the scheduling scope's THROTTLING ACCOUNTING — on cgroup v2, `cpu.stat` `nr_throttled` / `throttled_usec` as DELTAS over a measured window under load, as a fraction of enforcement periods, WITH a quiet-phase control — NEVER from averages: quota enforcement is per-period, so a workload far below quota ON AVERAGE can still throttle a large fraction of periods and stall every co-resident process, while CPU average / load average / RSS / thread counts all read HEALTHY (their quiet zeros are §11.4.201(6) FALSE-NULLS; forensic FACT 2026-07-22: a session scope at quota 9.6 CPU/100ms on a 64-CPU host, average demand 4.35 — UNDER quota — yet 16.9% of periods throttled during active phases, 0% quiet, worsening with session age, the co-located terminal-multiplexer server freezing ≥100ms per throttled period into multi-second interactive stalls); an "averages look healthy" verdict without throttle-accounting deltas is a §11.4.6 guess, and the same discipline binds every quota'd scheduler (container CPU limits, VM caps/steal). DESIGN half: an interactive-latency-sensitive process

(terminal server / TUI / editor / operator shell) MUST NOT share a fixed-quota NO-BURST scope with a bursty agent/build fleet — grant a measured burst allowance, isolate the interactive process in its own scope, or both; quota values are HOST-ADAPTIVE (computed from detected capacity per §12.11, §11.4.6 never hardcode — sibling FACT: a FIXED 2-CPU default quota on a 64-core host → 18.8% of periods throttled, 1757s forced idle, fixed by host-adaptive computation proven with a four-quadrant polarity test reading the LIVE cgroup back per §11.4.108); the THIRD orthogonal host-exhaustion axis beside §12.6 memory + §12.12 threads — no axis's instrument substitutes for another's. Honest boundary §11.4.6 — throttle telemetry proves WHERE periods were lost, not that throttling is the only degrader (RTT floors / GC load ranked separately in the founding investigation), and the design half never licenses removing a quota host policy requires (§12).

Classification: universal (§11.4.17). Composes §11.4.6/.24/.102/.108/.128(1)/.201/§12.6/§12.8/§12.11/§12.12. Propagation gate CM-COVENANT-114-225-

PROPAGATION (literal 11.4.225) + recommended gates CM-QUOTA-THROTTLE-TELEMETRY + CM-INTERACTIVE-SCOPE-NOT-QUOTA-STARVED + paired §1.1

mutations (averages-only verdict on a progressive-degradation issue → telemetry gate FAILs; fixed low quota literal on a detected many-core host, or terminal server co-located with the fleet in a no-burst quota scope → isolation gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item). **Canonical**

authority: constitution submodule Constitution.md §11.4.225. Non-compliance is a release blocker. No escape hatch — no --averages-prove-healthy, --skip-throttle-telemetry, --fixed-quota-default-OK, --interactive-in-fleet-scope-OK, --no-burst-needed, --skip-quiet-phase-control flag.

EXTENSIONS landed 2026-07-23 (the systematic-debugging / measurement-semantics harvest — full text in Constitution.md ; every lesson verified against captured evidence or re-demonstrated under control conditions before codifying, §11.4.6). §11.4.115(G) — defect-identity binding for

CONSTRUCTED reproductions: a RED test's PRECONDITION must be TRACEABLE to the reported defect's real evidence (the reporter's sequence §11.4.199, a captured failure trace, or a probe proving the precondition occurs in the reported environment), recorded with the RED verdict as a precondition-provenance field { observed | constructed }; when NO working repro exists and the precondition is CONSTRUCTED, the fix mints at most DEFENSIVE HARDENING — labelled as such, NEVER closing the reported item (it stays OPEN §11.4.197); systematic-debugging refutation of every mechanism a pending fix

addresses REVOKES its done-claim on the spot (forensic FACT 2026-07-22: an exit-hang fix nearly shipped on a hand-made-lock-holder RED → GREEN; real-PTY §11.4.199-shape re-driving refuted every addressed mechanism — 6/6 clean exits — and the change was honestly re-labelled in-source "defensive hardening ... NOT a fix for the reported exit-hang"); strengthens (A) from artifact-synthetic to MECHANISM-synthetic; gate `CM-RED-PRECONDITION-DEFECT-TRACEABLE` (a constructed provenance on a defect-CLOSING claim FAILs; the same provenance on an honest hardening item PASSes — the §11.4.201(1) false-positive guard) + paired §1.1 mutation. **§11.4.201(9)(10)(11)(12)** — measurement-semantics quartet: (9) FIELD-IDENTITY + UNITS + SEMANTIC CLASS {capacity|occupancy|rate|cumulative|high-water} verified against the authoritative source definition before ANY summary-read metric gates or propagates — CAPACITY≠OCCUPANCY is the named canonical instance (forensic FACT: `rcv_wnd 181120`, a window-CAPACITY ceiling, misread as a "~1.4s standing send queue"; real Send-Q a few hundred bytes `app_limited` ; propagated into THREE agent briefs + operator reports before the raw field was re-read, then publicly withdrawn — a downstream consumer of an unverified metric INHERITS the bluff); (10) OBSERVER DECONTAMINATION — an instrument contributing to its own measured population (probe sockets in a socket census, sampler load in a load profile) identifies + EXCLUDES its footprint AND re-proves the decontaminated instrument still sees via a (7)(b) control needle (FACT: TIME-WAIT 68–80 raw → 4–5 decontaminated, ~16× observer-induced, while the same scan still caught a real +264MB/1 → 93-socket step); composes §11.4.128(1) observer-effect budget + §11.4.119 single-resource-owner for measurement arms; (11) ARTIFACT-USABILITY-NOT-PREREQUISITE-PRESENCE — a readiness/install/health gate probes the ARTIFACT through its real invocation path (§11.4.224(A)), never a build prerequisite's presence (FACT: a `command -v <compiler>` gate turned a host holding a WORKING prebuilt into a false "cannot launch" exit-1 — "[ok] verified" and "[FAILED] cannot launch" in ONE transcript, review-graded BLOCKING; the proxy fails BOTH directions — clause (1) and (2) at once); (12) the SHELL-INSTRUMENT FOOTGUN CHECKLIST — NEW standing reference `docs/guides/shell_instrument_footguns.md` (inherited by reference §11.4.28/§11.4.177): exec-wrapper classes (`timeout / nice / env`) cannot run shell functions/aliases (`rc=127` false refusal), `pgrep -f self/carrier` match (§11.4.196(D)/§12.12), delimiter-keyed extractors TRUNCATING the unit under test into a FALSE PASS, `bash -x xtrace` BLINDED by the code under test's

`exec ... 2>/dev/null (countermeasure BASH_XTRACEFD )`, and a background watchdog subshell spawned INSIDE a `$(...)` command-substitution — it inherits the substitution's pipe WRITE-END, and an early disarm (`kill`) ORPHANS its `sleep` grandchild which keeps that write-end OPEN, so the `$(...)` BLOCKS for the FULL timeout budget after the real work completed instantly, every verdict and exit code CORRECT (countermeasure: redirect the watchdog subshell's fds away from the pipe — `( ... ) >/dev/null 2>&1 &` — or have the probe write to a file) — all five DEMONSTRATED under control conditions 2026-07-23; every shell-touching §11.4.142/§11.4.209 review consults it, every shell-instrument zero stays bound by the (7)(b) class-matched needle; CM-GUARD-ASSERTS-REAL-CONDITION EXTENDED with invariants (iv) metric-semantics-verified / (v) census-decontaminated / (vi) artifact-probed-not-prerequisite / (vii) checklist-consulted + paired §1.1 mutations (gate-code = separate work item). **§11.4.67(6)** — sourced-function `exec` redirection-scoping: a command-less `exec` applies EVERY redirection on its line to the CURRENT shell PERMANENTLY, so in sourced/library functions an fd-manipulation `exec` MUST brace-scope its side redirections (`{ exec 9>>"$f"; } 2>/dev/null`); a bare `exec N>file ... 2>/dev/null` in sourced code is a review defect + greppable gate class (forensic FACT 2026-07-22: a sourced lock helper's bare `exec 9>>lock 2>/dev/null` silenced the operator's interactive-shell stderr for the shell's LIFETIME — one provider launch left the terminal unable to print any error — while the paired unlock had used the correct brace form all along; marker-probe proven pre/post-fix); gate CM-SHELL-EXEC-REDIRECTION-SCOPED (brace-scoped forms and `exec -with-command` exempt — the §11.4.201(1) false-positive guard) + paired §1.1 mutation. **§11.4.142** — teaching forensic FACT added (2026-07-22): a transport-hardening change was committed AND pushed to `main` BEFORE any independent review (honestly flagged); the post-hoc §11.4.209-substrate review returned NO-GO with ONE BLOCKING finding — an unsafe security-relevant default, already public on every mirror — plus 2 Important + 4 Minor, and remediation consumed a full fix-forward cycle a pre-push review would have folded into the original change; the rule needed NO change (§11.4.142/.125/.134/.209 already bind it with no carve-out) — the incident is cited as its newest forensic anchor; remediation fix-forward per §11.4.113, never a history rewrite. ALREADY-COVERED verdicts this round: measurement-arm single-resource-ownership + observer-effect budget (§11.4.119 + §11.4.128(1) — the decontamination instance landed as §11.4.201(10), not a new anchor); review-before-push as a rule (bound in full by §11.4.142 family — teaching

FACT only). All Classification: universal (§11.4.17) — every toolchain literal genericised; consumers supply scopes, quotas, instruments, and paths as DATA per §11.4.35. The literals 11.4.115 / 11.4.201 / 11.4.67 / 11.4.142 are unchanged — their propagation gates stay GREEN.

§11.4.226 — Evidence-class-at-closure + standing detection pressure: machinery presence does NOT predict whether a fix holds — the EVIDENCE CLASS at closure (runtime vs source) under real DETECTION PRESSURE does; prose does not bind, seams do (research-derived, 2026-07-23). Compact summary: the decisive finding of the 2026-07-22 reopen/first-touch root-cause forensics, confirmed across four independent evidence lines (internal forensics; a 76-source external literature pass; one cross-project corpus WITHOUT tracking machinery; one WITH exemplary machinery): **every fix whose done-claim was preceded by an on-target RED → GREEN polarity flip that held recorded ZERO reopens; every fix confirmed only at source-green bounced** (runtime-evidence corpus: median attempts-to-stick 1, ~0.3% corrected reopen — a floor under weak detection pressure; source-green corpus: 52% recorded, itself an undercount; the healthy corpus's ONLY bouncing family was a wrong-layer echo-assertion; 69 PENDING-BUILD claim-moves; the last audited release entered manual QA with 1/25 items user-layer-validated — the operator was the FIRST user-layer oracle most claims ever met, so first-touch rediscovery was the deterministic pipeline output, not bad luck). The mandate — the STRICT FORMALIZATION of §11.4.115(F)/§11.4.146(D3)/§11.4.108: (1) closed evidence-class RANKING {runtime > artifact > source} with a per-defect-layer floor at every closure seam — a user-visible/runtime defect can NEVER close on artifact/source-class evidence (source-on-source stays legal, the §11.4.201(1) false-refusal guard); (2) class proven by MACHINE FIELDS never a label (runtime ⇒ target fingerprint read FROM the target + runtime observable; artifact ⇒ path + content hash; source ⇒ ref; label-without-fields = SHAPE-INCOMPLETE, unreadable = BLIND); (3) ANTI-ECHO — a "runtime observable" that IS a grep/source-scan transcript is WRONG-LAYER, refused (a grep can never observe a pixel; the pixel-defect-by-five-greps closure dies at this seam); (4) per-status evidence PRECONDITIONS — every §11.4.15/§11.4.33 status carries a checked precondition, not only the terminal write (FACT: a penultimate "ready"-class status parked 79 items — 42% of the done-claiming population — as design-complete-but-UNVERIFIED; the operator's 6/6-broken sample was drawn from exactly that set); (5) FIX-KIND honesty — a state-only repair/mitigation cannot close a defect item as fixed (sibling-corpus FACT: 11

recurrence chains, median 3 / max 7 attempts, dominated by state-only "fixes"); (6) STANDING DETECTION PRESSURE — pressure is a CAUSAL INPUT to every measured quality rate (a rate without its pressure regime is incomparable, §11.4.6; the 52%-vs-~0.3% split is explained by pressure + closure class, not machinery; the healthy corpus's one recurrence surfaced at 24 days = the next SWEEP not the next use); every registered, topology-present guard carries a FRESHNESS CONTRACT (verdict on the CURRENT artifact fingerprint within a declared staleness budget); the never-executed/stale/over-budget set feeds a STANDING risk-ordered re-run queue (most-reopened-first §11.4.189, then stalest-first) drained by the zero-idle loop (§11.4.94/§11.4.97) as regular development-cycle work — the §11.4.135 release seam is the blocking BACKSTOP, never the only executor (FACT: two standing guards, finally executed during the forensics, immediately emitted the latent FAILs — one reproduced the operator's loudest defect 3/3; registration was treated as coverage while execution was nobody's job); topology-absent guards enumerated never queued, an empty registry is BLIND never "all fresh". Binding: applies during EVERY regular development cycle; CHECKED through every code-review per §11.4.194(6); on every constitution pull the §11.4.32 sweep + §11.4.164 hook analyze + apply the additions and violations are tracked + cleared (§11.4.15) — no parallel mechanism. Honest boundary §11.4.6: a measured correlation with a stated mechanism and zero counter-cases — NOT a randomized trial (partially-censored exposure carried openly); fabricated machine fields stay closed by §11.4.115(F) harness-written verdicts + the §11.4.135 candidate-fingerprint join (layered, not eliminated); oracle calibration stays §11.4.107(10); the layer taxonomy is authored data (mislabelling is the §11.4.194(6) (a) review class's detective target); re-running a weak guard proves little — oracle strength stays §1.1/§11.4.115(F), orthogonal discovery stays §11.4.85/§11.4.118. Reference mechanisms (proven, NOT wired §11.4.205): SOL-01/02/04/09 RED-first POCs with golden-good/golden-bad/negative-control fixtures under docs/research/quality/solutions/ (65/65 GREEN, deterministic second iteration) — no force until their seams are wired by tracked §11.4.197 work items. Classification: universal (§11.4.17) — consumers supply layer bindings, evidence stores, verdict registries, and staleness budgets as DATA per §11.4.35. Propagation gate CM-COVENANT-114-226-PROPAGATION (literal 11.4.226) + recommended gates CM-EVIDENCE-CLASS-AT-CLOSURE + CM-GUARD-FRESHNESS-SCHEDULER + paired §1.1 mutations (close a runtime-layer item on source evidence / label-without-fields / echo-observable / stale guard with no queue entry → the gates FAIL; strip the literal → propagation FAILs; gate-code = separate work item, NOT claimed shipped

§11.4.6). **Canonical authority:** constitution submodule `Constitution.md`

§11.4.226. Non-compliance is a release blocker. No escape hatch — no

`--source-green-closes-runtime-defect` , `--label-is-class` , `--echo-observable-OK` , `--registration-is-coverage` , `--rate-without-pressure-regime` , `--mitigation-closes-defect` , `--skip-freshness-queue` flag.

§11.4.227 — Governance-corpus self-custody: the rule corpus is bound by the same custody it imposes — every named gate implemented-or-registered-deferral under a monotone ratchet, and propagation gates count anchor BLOCKS (exactly-once + lockstep-identical), never bare literals (research-derived, 2026-07-23). Compact summary: measured on THIS corpus 2026-07-22 (commands + control needles preserved in the root-cause analysis): 10,735 canonical lines, 220 distinct anchors; **413 named CM-* gates — 241 (58%) with no implementation in either canonical gate site — plus 85 literal "gate-code = separate work item" deferrals**; the custody seam designed 2026-07-17 was still prose on 2026-07-22 (the corpus's highest-value remediation fell into its own gap within five days); the fix-confirmation anchors predate the failures they failed to prevent by five-plus weeks — while **every fully-mechanised rule in the record either prevented or exposed its failure mode** (the PENDING-semantics + release verdict-coverage seams landed within a day of diagnosis and now block): the gap is not rule quality, it is the prose-to-seam conversion rate. The SHAPE defect: the propagation-gate family — the only family that runs everywhere — asserts literal-presence ≥ 1 , mechanically rewarding restatement over implementation ("the corpus measures its own health in words present, and so it grows words"); consequences measured: two anchors duplicated verbatim-divergently in ALL FOUR mirrors while every presence-gate stayed green (the F7 incident); one anchor NUMBER heading two different mandates (the §11.4.140/ §11.4.141 collision, independently re-detected by the block-integrity instrument — resolution stays an operator-owned re-mint); two canonical entry-point mandates contradicting (the mandated resumption file unreachable from the mandated README); external corroboration: mandated-checklist adoption without workflow embedding moved nothing across 101 hospitals (the Ontario null). The mandate: **(A) NAMED-GATE LEDGER + RATCHET** — every `CM-*` -class token named in the corpus is, at every build, either IMPLEMENTED (token in an EXECUTABLE gate site; prose carriers never count §11.4.201(7)(a)) or covered by a REGISTERED DEFERRAL row pointing at a tracked §11.4.197 item; silent gate debt refused; the unimplemented count is MONOTONE-DECREASING (validated

against the definition of done per §11.4.201(8) — at the correct end-state it is 0); a vanished gate NAME without an explicit removal citation FAILs (the delete-names-to-lower-the-count gaming channel closed; repeals stay legal + visible, the §11.4.166 pattern); **an anchor's "done" state is its SEAM landing, not its TEXT landing** — a rule authored without its seam is in §11.4.205's worse-than-nothing state, and the ledger makes that state visible, counted, and shrinking; brownfield baseline = operator §11.4.66 decision (the §11.4.224(E) adoption-fence pattern), never an invented ratchet. **(B) ANCHOR-BLOCK INTEGRITY** — every propagation-class gate counts BLOCK-STARTS (structural line-anchored heading patterns), never bare literals: exactly ONE block per anchor per governance file (0 = absence, >1 = DUPLICATION, FAIL naming anchor + file); content-hash EQUALITY across the declared §11.4.157 lockstep mirror set (divergent copies FAIL naming both files; compact-vs-full variance ACROSS layers stays legitimate — canonical-vs-mirror drift is the §11.4.186 family's job); mid-body citations are CARRIERS and never count; ZERO blocks extracted = BLIND-OR-EMPTY never clean (§11.4.201(6)); an anchor NUMBER heading two mandates is a COLLISION and FAILs (§11.4.54 never-reuse applied to anchor ids); the extractor captures full dotted ids so a sub-anchor never prefix-matches its parent (a measured §11.4.201(1) false positive in the founding instrument, fixed RED-first). **(C)** semantic contradiction between two DIFFERENT anchors stays undecidable-by-hashing — only the decidable projections are mechanized (fenced scopes §11.4.112(5), seam declarations §11.4.120), the rest is §11.4.142/§11.4.194 review territory, stated never claimed; this anchor adds the enforcement-ratio INSTRUMENT, not a growth ban. Binding: the ledger regenerates + diffs on every governance edit + every build (regular development cycles, never an on-demand audit); review-checked per §11.4.194(6)(c); on every constitution pull the §11.4.32 sweep + §11.4.164 hook run the ledger + block-integrity checks and violations are tracked + cleared. Reference mechanisms (proven, NOT wired §11.4.205): SOL-05 + SOL-06 RED-first POCs under docs/research/quality/solutions/ (SOL-05's live run reproduced the 413-gate census exactly; SOL-06's live run confirmed the four mirrors clean post-F7 + re-detected the 140/141 collision) — no force until wired by tracked work items. Classification: universal (§11.4.17) — consumers (and this governance repo itself) supply gate-site lists, deferral registries, and lockstep mirror sets as DATA per §11.4.35. Propagation gate CM-COVENANT-114-227-PROPAGATION (literal 11.4.227 — itself REQUIRED block-start-shaped per (B)) + recommended gates CM-GATE-LEDGER-RATCHET + CM-ANCHOR-BLOCK-INTEGRITY + paired §1.1 mutations (name a gate with no implementation + no deferral / delete a name

without citation / duplicate a block / diverge two lockstep copies / re-mint a number → the gates FAIL; strip the literal → propagation FAILs; gate-code = separate work item, NOT claimed shipped §11.4.6). **Canonical authority:** constitution submodule `Constitution.md` §11.4.227. Non-compliance is a release blocker. No escape hatch — no `--name-gate-without-seam` , `--silent-gate-debt-OK` , `--delete-name-to-lower-count` , `--presence-gte-1-suffices` , `--duplicate-anchor-OK` , `--mirror-drift-OK` , `--reuse-anchor-number` flag.

§11.4.228 — Cross-agent extension lifecycle mandate: per-platform compatibility, source tracking, documentation with compatibility matrix, and auto-wiring on constitution pull (research-derived from Extensions Catalog forensics, 2026-07-24). Compact summary: every governed extension (Skill, MCP, ACP, LSP, plugin) MUST (A) ship a per-platform compatibility declaration (Claude Code, OpenCode, Gemini CLI, Qwen Code minimum) with closed vocabulary `{SUPPORTED|UNTESTED|UNSUPPORTED|CONFIG-ONLY}` , NEVER cause loops/broken behavior in any platform, and carry a per-platform liveness test §11.4.169; (B) be traceable to its exact canonical origin (repo + commit SHA + path, recorded in `EXTENSION_SOURCE.yaml` — symlinks preferred over copies, blind copies without provenance are §11.4.124 dead-code-in-waiting), with a per-project extensions catalog (`docs/extensions/EXTENSIONS_CATALOG.md`) as the single source of truth, §11.4.86 roster/corpus-backed fingerprint; (C) the extensions catalog MUST carry a per-platform compatibility matrix + liveliness-testing evidence per extension × platform per §11.4.69 `feature_class=extension_loading` , four-format export per §11.4.65, auto-synced via §11.4.106 docs-chain; (D) the §11.4.164 `post_update_hook.sh` is EXTENDED to auto-detect all governed agent platforms, symlink constitution skills into each (`.claude/skills/` , `.opencode/skills/` , platform equivalents), register MCP servers into each platform's config, self-validate, and be idempotent — closing the gap where 4 constitution skills exist in source but are un-wired to any platform (the §11.4.108 source-present-runtime-absent gap). Honest boundary §11.4.6 — wiring guarantees REACHABILITY, never runtime acceptance. Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-228-PROPAGATION` + recommended gates `CM-EXTENSION-PER-PLATFORM-COMPAT` / `CM-EXTENSION-SOURCE-TRACKED` / `CM-EXTENSION-CATALOG-DOCUMENTED` / `CM-CONSTITUTION-EXTENSIONS-AUTO-WIRED` + paired §1.1 mutations. Full text in `Constitution.md` §11.4.228.

§11.4.229 — Live in-session task/todo tracker MUST always be up to date and fully in sync with the real work state — never stale, never completed-shown-as-pending nor pending-shown-as-done (User mandate, 2026-07-25). Verbatim operator mandate: "make sure it is ALWAYS up to date fully in sync! CRITICAL: This MUST BE one of mandatory rules / constraints! Add it with all relevant details into the constitution Submodule and relevant related files." The agent's LIVE IN-SESSION TASK/TODO TRACKER — the harness-provided task list /

`TodoWrite` -class todo tool / any equivalent moment-to-moment execution-state surface the runtime exposes DURING a session — MUST ALWAYS be kept up to date and FULLY IN SYNC with the real live work state: never stale, never showing completed work as pending, never showing pending / blocked / abandoned work as done or in-progress; updated the MOMENT any tracked task's state changes (started → in-progress, completed → done, blocked, abandoned, split, re-scoped, newly discovered), never batched to a later checkpoint, never reconstructed from memory at the end, never left reflecting a superseded plan — a live tracker that lags the real work is a §11.4 PASS-bluff at the execution-transparency layer (an operator or resuming conductor per §11.4.116 makes decisions on a false picture of what is done / in-flight / owed). (A) every state-change updates it IMMEDIATELY (exactly one in-progress where the surface models that; no completed-marked-pending, no pending-marked-done, no ghost task for dropped work — dropped work is closed with a reason); (B) NEVER stale, NEVER inverted (§11.4.6 — the tracker is captured FACT, falsely-done = PASS-bluff / falsely-pending = §11.4.1 FAIL-bluff, drift corrected the moment observed); (C) DISTINCT artifact that COMPLEMENTS never REPLACES §12.10 (committed CONTINUATION doc) / §11.4.131 (standing committed session-resumption FILE) / §11.4.97 (operator-facing milestone progress cadence) / §11.4.106(F) (doc/DB sync at commit/build/constitution-pull write-seams) / §11.4.202 / §11.4.208 (workable-item DB + request-history ledgers) — those bind committed/cross-session/operator-report/tracked-work artefacts, NONE binds the live in-session tracker itself; it is the in-session analogue of §12.10's always-current mandate applied to the harness task list, the truth those artefacts are materialised FROM; (D) honest boundary §11.4.6 — where no such surface exists this is honestly N/A (§11.4.3 SKIP-with-reason, latent-binding per §11.4.96), and "in sync" = reflects the real work state as the agent knows it now, not omniscience about unobserved work. Classification: universal (§11.4.17). Composes §12.10 / §11.4.131 / §11.4.97 / §11.4.106(F) / §11.4.126 / §11.4.116 / §11.4.202 / §11.4.208 / §11.4.6 / §11.4.1. Propagation gate `CM-COVENANT-114-229-PROPAGATION` (literal `11.4.229`) + recommended gate `CM-`

LIVE-TASK-TRACKER-IN-SYNC + paired §1.1 mutation (gate-code = separate work item, NOT claimed shipped §11.4.6). No escape hatch — no `--tracker-may-lag` , `--batch-todo-updates` , `--stale-task-list-OK` , `--reconstruct-todos-at-end` , `--skip-task-tracker-sync` , `--pending-may-be-done` , `--completed-may-stay-pending` flag. Full text in `Constitution.md` §11.4.229.

§11.4.230 — Parallelized-pipeline methodology: build ↔ validate overlap + affected-stages-only re-run + parallel multi-device testing/recording + fan-out-subagents-always (research-derived, 2026-07-26).

Forensic motivation (genericised): the sequential pipeline shape — build FULLY completes → THEN validation begins → on ANY discovery, restart the WHOLE cycle from scratch → all devices tested one-at-a-time — serialises the operator's wall-clock onto build latency + one-device-at-a-time evidence gathering, so a multi-hour containerised build followed by serial per-device validation and a from-scratch restart-on-discovery makes time-to-a-manual-QA-capable-release a multiple of what the genuinely-parallel work requires. §11.4.58's PWU pipeline already overlaps Stage 1 of round N+1 with Stages 4–5 of round N; §11.4.89 already backgrounds long tests; §11.4.103 already mandates ≥3 auto-backfilled parallel streams; §11.4.119 already partitions shared hardware by single-owner; §11.4.128/§11.4.144 already mandate always-on recording + availability-following. §11.4.230 BINDS those into one operating discipline and adds the three invariants they did not each state on their own: (A) the build ↔ validate CONCURRENCY seam + affected-stages-only re-run-on-discovery, (B) all-devices-concurrent + worker-pool recording/processing, (C) fan-out-subagents as the standing default. It is the parallelization capstone the way §11.4.103 is the continuous-parallel-stream capstone — a strengthening bind, never a weakening.

(A) PIPELINE-OVERLAP PARALLELIZATION. Validation/verification work MUST run CONCURRENTLY WITH the build wherever a validation stage does NOT depend on the build's own not-yet-produced artifact — never deferred until the build finishes when it could have run alongside it. Deploy the moment the build is done (build-completion is the deploy trigger, per §11.4.121 no-commit-while-build-writes-tracked-artifacts + §11.4.108 artifact-must-exist ordering — the deploy waits for the artifact, the deploy-independent validation does not). Concretely, WHILE a build runs, the following run in parallel per the §11.4.96 SAFE-during-build catalogue: pre-build source-layer gates + meta-test mutation sweeps + coverage/clash checks (§11.4.110), the §11.4.125/§11.4.142/§11.4.194 independent code-review on the

batch's SOURCE, doc/DB sync (§11.4.106), and on-device validation of the PRIOR build (device-independent of the in-flight build, §11.4.119). On any NON-huge-blocker discovery, the cycle MUST: (1) bump the version code (a new, distinctly-identifiable candidate artifact — never re-use the failing artifact's id, §11.4.111 stable-name + §11.4.200 target-identity); (2) fix at the proven root cause (§11.4.102 systematic-debugging first, §11.4.115 RED-on-the-broken-artifact); (3) RE-RUN ONLY THE AFFECTED STAGES, PARALLELIZED — the stages whose inputs the fix changed (its own targeted validation FIRST per §11.4.130, then the affected regression set), NOT a serial restart-from-scratch of the entire cycle. The explicit efficiency target is a **4–5× reduction in time-to-a-manual-QA-capable-release** versus the fully-sequential shape — a measured target, tracked as build-and-validate wall-clock deltas (§11.4.24 build-resource stats), never an assumed one (§11.4.6). **HONEST BOUNDARY (§11.4.6 / §11.4.108 / §11.4.120)**: the CURRENT build's ARTIFACT-layer (§11.4.108 layer 2) and RUNTIME-ON-CLEAN-TARGET-layer (layer 3) validation CANNOT precede its own artifact — a gate whose precondition is the build's output belongs at a post-build seam, never overlapped onto authorship (§11.4.120 seam-placement; §11.4.69

`artifact_not_yet_built`). "Concurrent with the build" therefore means the deploy-INDEPENDENT stages overlap; the CURRENT-artifact runtime validation begins the moment that artifact exists, in parallel across devices per (B). **"Affected-stages-only re-run" is bounded to NON-huge-blocker discovery**: a HUGE BLOCKER (release-blocking-severity defect / regression invalidating the cycle / blast radius reaching the batch's other fixes) still triggers the §11.4.129 STOP → fix-all → full-RESTART protocol (restart, never resume) — §11.4.230(A) does NOT weaken §11.4.129, §11.4.40 full-suite-retest-before-tag, §11.4.130 validate-the-fix-first, §11.4.132 risk-ordered validation, or §11.4.185 manual-QA-final; it narrows the from-scratch restart to exactly the cases those anchors require it, and parallelises everything else.

(B) PARALLEL MULTI-DEVICE TESTING + RECORDING/PROCESSING. Every available, reachable test device MUST be exercised CONCURRENTLY — one validation stream per device, never one-device-at-a-time when the devices are independent — each stream the SINGLE EXCLUSIVE OWNER of that device's exclusive resources per §11.4.119 (the owner drives; any concurrent stream targeting the same device is read-only), so per-device evidence is never cross-contaminated (§11.4.119 evidence-integrity). Every device is ALWAYS-RECORDED per §11.4.128 (non-invasive, side-effect-free, deterministic layout) and

AVAILABILITY-FOLLOWED per §11.4.144 (a dropped device logs an honest offline event + reconnect-and-resume, never a silent corpus hole, never a faked continuity). Recording CAPTURE + ANALYSIS/PROCESSING MUST be fanned out through a BOUNDED WORKER-POOL (a Go worker-pool or platform equivalent) so N devices' recordings are captured + analysed in parallel and per-device evidence never serialises onto a single processor — the pool bounded by §12.6 memory + §12.12 thread headroom. Per-device PASS cites its own captured-evidence artefact (§11.4.5/§11.4.69/§11.4.107) on a CLEAN/fresh deployment (§11.4.108 runtime-signature, §11.4.200 verify-deployed-artifact-identity-per-device); the highest-risk / most-reopened cases (§11.4.132/§11.4.189) get the deepest scrutiny FIRST on every device. **HONEST BOUNDARY (§11.4.6):** genuinely-shared exclusive resources (one HDMI/audio sink, one media-playback path) serialise per §11.4.119 single-owner-over-time even across devices; "all devices concurrent" means independent devices, never two owners of one exclusive resource.

(C) FAN-OUT-SUBAGENTS-ALWAYS. Subagent fan-out (§11.4.20/§11.4.70 subagent-driven; §11.4.187 multi-track ruler; §11.4.192 auto-backfill) is the STANDING DEFAULT for EVERY genuinely-parallelizable effort — batch investigation, test-authoring, gate-authoring, per-device validation (B), per-PWU development (§11.4.58), doc/research work — not a per-request opt-in and not reserved for "big" tasks. A parallelizable effort executed monolithically in the foreground when a fan-out path exists is a §11.4.230(C) violation (the §11.4.94/§11.4.97/§11.4.103 zero-idle discipline applied to HOW work is dispatched, not only WHETHER it runs). **HONEST BOUNDARY (§11.4.6 / §11.4.183):** "always fan out" is BOUNDED by (1) genuine parallelizability (contending/redundant/serially-dependent subagents are NOT fan-out — they are waste, and spawning them is NOT maximization per §11.4.183 maximum-USEFUL-throughput); (2) §12.6 60% memory ceiling; (3) §12.12 `RLIMIT_NPROC` thread headroom; (4) §11.4.58 ≤ 6 concurrent-agent cap; (5) §11.4.119 single-resource-owner (two subagents may not both drive one exclusive device); (6) critical-state sequencing (commits/pushes/tags stay serial per §11.4.88/§11.4.113). The bar is maximum USEFUL parallel throughput, never agent count; fan-out that exceeds a host-safety ceiling or contends a single resource is a violation, not compliance.

Classification: universal (§11.4.17) — a platform-neutral parallelization discipline reusable by ANY project that builds an artifact, validates it across multiple targets, and records device evidence; the consuming project supplies its concrete device

set + serials, build/flash tooling, recording layout, and worker-pool implementation per §11.4.35. Composes §11.4.20 / §11.4.24 (wall-clock stats prove the 4–5× target) / §11.4.40 / §11.4.58 (the PWU pipeline this binds + extends) / §11.4.69 / §11.4.70 / §11.4.88 / §11.4.89 / §11.4.94 / §11.4.96 (SAFE-during-build catalogue) / §11.4.97 / §11.4.102 / §11.4.103 / §11.4.108 (artifact/runtime layer ordering — the honest boundary) / §11.4.110 / §11.4.115 / §11.4.119 (single-resource-owner — B's load-bearing partition) / §11.4.120 (seam-placement — A's honest boundary) / §11.4.121 / §11.4.125 / §11.4.126 / §11.4.128 (always-on recording — B's evidence source) / §11.4.129 (huge-blocker full-restart — the un-weakened carve-out) / §11.4.130 / §11.4.132 / §11.4.144 (availability-following — B's drop-handling) / §11.4.183 (maximum-USEFUL-parallelism — C's honest boundary) / §11.4.185 (manual-QA-final — the un-weakened terminal gate) / §11.4.187 / §11.4.192 / §11.4.200 (per-device deployed-artifact verification) / §12.6 / §12.12. Propagation gate `CM-COVENANT-114-230-PROPAGATION` (literal `11.4.230` across the consumer fleet) + recommended gates `CM-PIPELINE-OVERLAP-VALIDATION` (a build that runs with zero deploy-independent validation stages running concurrently when SAFE-catalogue work is queued → FAIL; a NON-huge-blocker discovery that restarts the whole cycle from scratch instead of re-running only the affected stages → FAIL) + `CM-PARALLEL-DEVICE-WORKER-POOL` (available independent devices validated serially instead of concurrently, OR recording capture/analysis serialised onto one processor when a worker-pool path exists → FAIL) + `CM-FAN-OUT-SUBAGENTS-DEFAULT` (a genuinely-parallelizable effort run monolithically in the foreground when a fan-out path exists AND all host-safety ceilings have headroom → FAIL) + paired §1.1 meta-test mutations (per §1.1; gate-code = separate work item, NOT claimed shipped §11.4.6).

Canonical authority: constitution submodule `Constitution.md` §11.4.230.

Non-compliance is a release blocker regardless of context. No escape hatch — no

```
--sequential-build-then-validate ,
--restart-whole-cycle-on-discovery , --serial-device-testing , --
serialise-recording-processing , --no-fan-out , --monolithic-
parallelizable-work flag exists.
```

§11.4.231 — Nano-precision model-tier selection: lightest-capable-model-first with dynamic mid-task escalation (User mandate, 2026-07-26, CRITICAL).

Verbatim operator mandate: every work effort MUST use the LIGHTEST CAPABLE model, escalating tier only as complexity rises — a nano-precision, per-task model-tier choice, NOT "everything on the heaviest model." Forensic motivation (genericised, §11.4.6): dispatching every task — a trivial doc-sync, a one-line status read, a mechanical format fix — on the heaviest available model wastes capacity + token budget + wall-clock that a lighter tier would spend identically well, while the OPPOSITE failure (a hard multi-factor root-cause investigation dispatched on a tier too light to reason through it) produces exactly the shallow, guessed, or wrong-conclusion work §11.4.6 / §11.4.102 / §11.4.194 forbid. Neither "always heaviest" nor "always lightest" is compliant; the correct discipline is a PER-TASK, EVIDENCE-BASED tier choice that tracks the task's own complexity, re-evaluated as that complexity becomes known — the nano-precision selection this anchor mandates.

(A) LIGHTEST-CAPABLE-FIRST + ESCALATE-ON-COMPLEXITY. Every dispatched work effort — main-stream or subagent (§11.4.20/§11.4.70), single-track or multitrack (§11.4.187/§11.4.192) — is assigned the LIGHTEST model tier from the closed, capability-ordered ladder `haiku < sonnet < opus` genuinely sufficient to complete it correctly (the review/merge-conflict-resolution substrate is a SEPARATE, already-pinned tier per clause (E), not a fourth rung of this ladder). Tier selection is a DECIDABLE heuristic, not a guess (§11.4.6) — evaluated against, at minimum: (i) **task scope** (single mechanical edit / focused single-file change / multi-file-or-multi-component change); (ii) **reasoning depth** (a lookup-and-apply action vs. a multi-factor root-cause investigation per §11.4.102/§11.4.194); (iii) **blast radius** (a change touching one isolated file vs. one crossing shared state, hardware safety §11.4.133, or governance/gate seams §11.4.227); (iv) **novelty** (a previously-solved, pattern-matched task vs. a genuinely new problem with no established precedent). Closed-set tier bindings (illustrative, non-exhaustive — the consuming project supplies its own precise mapping per §11.4.35): `haiku` — trivial/mechanical work (simple greps/searches, doc-sync regeneration, status/field reads, format/lint fixes, boilerplate scaffolding); `sonnet` — moderate work (a focused single-component debug, a bounded single-file implementation, straightforward research or audit reads); `opus` — complex work (architecture/design decisions, multi-factor reasoning per §11.4.194(1), cross-feature blast-radius analysis per §11.4.92 Pass 2/3, hard root-cause investigation per §11.4.102 that resists a first-

pass hypothesis). A task whose complexity is genuinely UNDECIDABLE from available evidence starts at the NEXT tier UP from the best available estimate (the safe-reversible default per §11.4.101) rather than guessed downward.

(B) DYNAMIC MID-FLIGHT SWITCHING. Tier assignment is NOT a one-time, start-of-task decision — it is RE-ASSESSED continuously as the task's real complexity becomes known during execution. A task DISPATCHED as simple that PROVES harder mid-flight (an unexpected multi-factor interaction surfaces, a root cause resists the first two hypotheses, the blast radius turns out to cross a governance/hardware-safety seam) MUST ESCALATE to the next-heavier tier for its remainder — continuing on an under-powered tier past the point complexity is KNOWN is a clause-(D) false economy, not compliance. Conversely, a task DISPATCHED as heavy that DECOMPOSES into genuinely light, bounded, low-blast-radius subtasks (mechanical file edits following an already-designed pattern, doc regeneration, per-item status reads) dispatches THOSE subtasks on a LIGHTER tier via subagent fan-out (§11.4.20/§11.4.70/§11.4.183) rather than running every subtask on the parent's heavier tier by default — capability-matching applies at EVERY dispatch granularity, not only the top-level task.

(C) ALWAYS-ON COMPOSITION, NEVER A PER-REQUEST OPT-IN. Nano-precision tier selection is a STANDING DEFAULT mechanism, engaged automatically for EVERY work effort from the first prompt of a session onward (mirrors §11.4.126 default-autonomous-loop / §11.4.198 default-mechanisms-always-on) — applied TOGETHER WITH, never as a substitute for: §11.4.141 (heavy token-optimization — a lighter model dispatched with a bloated context is not the efficiency this anchor targets); §11.4.183 (maximum-USEFUL-parallelism — tier-matching each parallel stream, not merely running more streams on one fixed tier); §11.4.187/§11.4.192 (the multi-track ruler orchestration + continuous auto-backfill — each track's worker is tier-assigned per this anchor, not defaulted to one blanket tier project-wide); and §11.4.196/§11.4.198 (native-alias-first priority + always-on default mechanisms — this anchor governs WHICH TIER within whichever alias/account is selected, composing with, never replacing, §11.4.196's ALIAS-layer priority). None of these five composed anchors is weakened, duplicated, or re-specified by §11.4.231 — this anchor adds the TIER axis to their existing ALIAS/PARALLELISM/TOKEN axes.

(D) HONEST BOUNDARY (§11.4.6) — "LIGHTEST CAPABLE" MEANS

CAPABLE, NOT CHEAPEST-REGARDLESS. "Lightest capable" is a CONJUNCTION, not a discount: a tier choice that fails to actually complete the task correctly — producing a shallow investigation, a guessed root cause, an under-verified fix, or ANY form of bluff per the §11.4 anti-bluff covenant — is NOT "lightest capable," it is UNDER-POWERED, and under-powering a task that then fails, bluffs, or requires costly rework is a FALSE ECONOMY and a §11.4 violation of equal severity to any other PASS-bluff (the token/time saved at dispatch is spent many times over in re-investigation, re-review, and the risk of a shipped defect). Capability determinations are made from the clause-(A) decidable heuristic + observed task outcomes — EVIDENCE-BASED, never guessed (§11.4.6): "this always works fine on haiku " without a captured basis is exactly the guess §11.4.6 forbids, and a tier's genuine capability for a task class is established by the same anti-bluff discipline (§11.4.5/§11.4.69/§11.4.107) that governs every other claim in this constitution. A platform-chosen serving model outside the haiku < sonnet < opus ladder (e.g. fable served for general non-review work) is NOT a §11.4.231 dispatch-tier violation — the ladder binds the DISPATCHER's tier ASSIGNMENT, not the platform's own per-turn serving choice.

(E) HARD PINS OVERRIDE THIS DEFAULT FOR THEIR SCOPES. §11.4.231 sets the DEFAULT working tier for NON-review, non-merge-conflict-resolution work. It does NOT relax, override, or compete with the two existing hard model-substrate pins: §11.4.209 (every mandatory independent code-review — §11.4.125/§11.4.142/§11.4.194 iterated to GO per §11.4.134 — runs on **Fable at xhigh effort, Opus at xhigh** on genuine Fable-unavailability, NEVER a lower tier or lower effort) and §11.4.211 (every main → feature/product/flavor merge-CONFLICT resolution runs on the SAME Fable-xhigh / Opus-xhigh substrate). Those two anchors are STRICTER, SCOPE-SPECIFIC pins that this anchor's general ladder does not reach — review and merge-conflict-resolution are NOT eligible for haiku/sonnet/opus-by-complexity dispatch; they run on their pinned substrate ALWAYS, regardless of how "simple" the reviewed change or the conflict appears. A review or merge-conflict resolution dispatched on any tier other than the §11.4.209/§11.4.211 pin (absent a genuinely-proven Fable-unavailability fallback) is a violation of THOSE anchors, not a valid application of §11.4.231's general ladder.

(F) EFFORT-TIER SELECTION — the SAME nano-precision discipline applied to the reasoning-EFFORT axis (operator IMPORTANT mandate, 2026-07-26; verbatim: "Besides dynamic determining proper models to use, we MUST

decide the effort as well and switch between the effort settings same as we do now for chosen model! Labeling MUST BE extended for this details as well. Example: (T1/main - claude4 - opus - high)."). Every dispatched work effort is assigned not only a model tier per clauses (A)-(E) but ALSO the LIGHTEST reasoning-EFFORT setting from the closed ladder

`low < medium < high < xhigh < max` genuinely sufficient to complete the task correctly — chosen by the SAME clause-(A) decidable heuristic (scope / reasoning-depth / blast-radius / novelty), NOT a guess (§11.4.6). Illustrative bindings (non-exhaustive; the consuming project supplies its precise mapping per §11.4.35):

`low` / `medium` — trivial / mechanical work (doc-sync regeneration, status / field reads, format / lint fixes, boilerplate); `high` — moderate-to-complex reasoning (focused single-component debug, multi-factor analysis per §11.4.102 / §11.4.194); `xhigh` / `max` — the HARDEST verification / merge-conflict / architecture work.

(F.1) DYNAMIC MID-TASK ESCALATION — effort, like model tier per clause (B), is RE-ASSESSED continuously as the task's real complexity becomes known; a task dispatched at a lower effort that PROVES harder mid-flight MUST ESCALATE its effort for the remainder (continuing under-effort past the point complexity is KNOWN is the clause-(D) false economy), and a heavy task decomposing into light bounded subtasks dispatches THOSE at a lighter effort — effort-matching applies at EVERY dispatch granularity, not only the top-level task. (F.2) HONEST HARNESS BOUNDARY (§11.4.6, load-bearing — do NOT overstate the mechanism) — effort-tiering is REAL where the execution path exposes an effort control and a DOCUMENTED CAPABILITY-GAP where it does not: the current Claude Code **Agent tool** (subagent dispatch) exposes a `model` parameter but **NO effort parameter**, so a subagent's effort is NOT presently settable via that path (the §11.4.182 `<effort>` label field degrades honestly to `?` there); the

Workflow tool's `agent()` DOES take an `effort` argument, so effort IS settable there — the label reports whichever path is live, never fakes one. This is the SAME aspirational-vs-harness state the clause-(E) §11.4.209 / §11.4.211 Fable- `xhigh` pins already occupy: the pins mandate the effort, and whether a dispatch path can SET it is a harness capability reported honestly, never faked (§11.4.182). (F.3) HARD-PIN FLOOR PRESERVED — clause (E)'s hard pins are ALSO effort pins: every §11.4.209 mandatory independent code-review AND every §11.4.211 merge-conflict resolution runs at `xhigh` effort (Fable, or Opus at `xhigh` on genuine Fable-unavailability), and this general effort ladder NEVER tiers those two scopes below `xhigh` regardless of how "simple" the reviewed change or the conflict

appears — a review / merge-conflict resolution run below `xhigh` is a §11.4.209 / §11.4.211 violation, not a valid application of this ladder. (F.4) HONEST BOUNDARY (§11.4.6) — "lightest capable effort" is a CONJUNCTION, not a discount: an effort setting that yields a shallow / guessed / under-verified result is UNDER-POWERED, a FALSE ECONOMY, and a §11.4 violation of PASS-bluff severity, exactly as an under-powered MODEL tier is per clause (D); capability determinations are evidence-based, never guessed (§11.4.6).

Classification: universal (§11.4.17) — a platform-neutral model-capacity-matching discipline reusable by any project dispatching work across a tiered model roster; the consuming project supplies its own precise tier ↔ task-class mapping and its own model-name ladder per §11.4.35 (the `haiku / sonnet / opus` ladder and the Fable/Opus-`xhigh` pins are this project's current concrete instantiation, not a universally-fixed vocabulary — a consumer with a differently-named tier roster maps its own ladder onto the SAME lightest-capable-first + dynamic-escalation + always-on-composition + honest-boundary + hard-pin-override structure).

Composes §11.4.6 (no-guessing — capability determinations are evidence-based) / §11.4.20 / §11.4.70 (subagent-driven dispatch — the granularity at which tier-matching applies) / §11.4.92 (blast-radius analysis feeds the clause-(A) heuristic) / §11.4.101 (safe-reversible default on undecidable complexity — escalate up, never guess down) / §11.4.102 / §11.4.194 (reasoning-depth signals for the heuristic) / §11.4.126 / §11.4.141 / §11.4.183 / §11.4.187 / §11.4.192 / §11.4.196 / §11.4.198 (the always-on composition set, clause (C)) / §11.4.209 / §11.4.211 (the hard-pinned review/merge-conflict substrate this anchor does NOT touch, clause (E)).

Propagation gate `CM-COVENANT-114-231-PROPAGATION` (literal `11.4.231` present as a block-start, exactly-once per governance file, lockstep content-hash equality across the consumer fleet per §11.4.227(B)) + recommended gate `CM-MODEL-TIER-LIGHTEST-CAPABLE` (a dispatched effort whose logged/declared model tier exceeds the clause-(A) heuristic's indicated tier for a task later confirmed trivial-and-successful-at-the-lighter-tier is a WARN/candidate finding — never a hard FAIL on tier choice alone, since tier-adequacy is only provable in retrospect per clause (D); the gate instead HARD-FAILS the two decidable violations: (i) a review or merge-conflict-resolution effort NOT run on the §11.4.209/§11.4.211 pinned substrate, and (ii) a task whose complexity was PROVEN mid-flight [an escalation event was logged, or a downstream failure/bluff was traced back to it] that did NOT escalate tier before continuing) + paired §1.1 mutation (gate-code = separate work item, NOT claimed shipped — §11.4.6; mutation plan: (1) dispatch a

code-review labeled as running on `sonnet` instead of `Fable-xhigh` → gate MUST FAIL; (2) inject a logged mid-task complexity-escalation event with NO corresponding tier-escalation record → gate MUST FAIL; (3) golden-FALSE fixture: a task that stays on its initially-assigned tier throughout, with no escalation event logged → gate MUST NOT fire, per §11.4.201(1) the false-positive guard).

Canonical authority: constitution submodule `Constitution.md` §11.4.231.

Non-compliance is a release blocker regardless of context. No escape hatch — no

```
--always-heaviest-model , --always-lightest-model ,  
--skip-complexity-reassessment , --fixed-tier-per-session , --relax-  
review-pin , --relax-merge-conflict-pin , --per-request-tier-opt-in
```

flag exists.

§11.4.232 — Anti-mess long-op orchestration: every long-op is registered, single-owned-per-purpose, liveness-PROVEN, handed-off-on-stop, safely-reaped, consistency-checked, and escape-bounded — the single source of truth for what is running, never inferred (research-derived, 2026-07-26).

Forensic anchor (genericised): in one cycle a verification sweep produced five distinct mess states — (1) multiple concurrent same-purpose sweeps ran unintentionally (no single-owner-per-purpose claim); (2) a sweep HUNG on a consumer-blocked-on-an-unclosed-producer pipeline and was undetected because liveness was inferred from process-alive (a §11.4.201(6) FALSE-NULL: a wedged op and a progressing op returned the identical "process alive" signal); (3) stopping the owning subagent SILENTLY KILLED its child long-op because ownership was implicit (no handoff-on-stop); (4) a stale/non-growing log was read as live progress (the same false-null, reader-side); (5) an un-completable tooling gate blocked a release for hours (no escape to a terminal state). Per §11.4.102 Phase 4.5 these are ONE architectural state-management flaw, not five bugs: **there was no single source of truth for what long-op is running, who owns it, its state, and its PROVEN liveness** — the §11.4.196(F) configured~~in-use~~ gap generalised to project long-ops, so every mess state was un-observable, and an un-observable bad state is a §11.4/§11.4.1 bluff at the orchestration layer.

The mandate — a **long-op orchestration layer** binding every long-op (any op expected to exceed the runtime background-execution threshold per §11.4.89 — verification sweep, meta-test sweep, build, flash/deploy, test/QA suite, dispatched subagent), ALL of which hold:

(A) SINGLE SOURCE OF TRUTH — the long-op lifecycle registry. Every long-op MUST register BEFORE it runs and reach a terminal state (`complete|failed|reaped|handoff|blocked-escape`) when it ends, in a durable (never-tmpfs, write-temp-then-rename, §11.4.116 append-only event stream + atomic snapshot) registry recording `{op_id, purpose_key, owner, pid|pgid, cwd, state, last_heartbeat, heartbeat_seq, progress_offset, log_path, verdict, evidence_path}` . Success is read from the op's terminal verdict/ `result` , NEVER its process exit code. A long-op running with no registry record — or a record with no live op — is a FALSE-READY drift state.

(B) SINGLE-OWNER-PER-PURPOSE. No two same- `purpose_key` long-ops run at once; acquisition is a probe-race-safe claim keyed on `purpose_key` (§11.4.119). A live claim forces exactly ONE of REUSE (attach) or explicit SUPERSEDE (graceful handoff of the running owner, logged), never a silent second run.

(C) LIVENESS TRUTH-SOURCE. Progress is PROVEN, never inferred from process-alive: every long-op emits a monotonic heartbeat and/or an advancing progress offset; a watchdog reads the DELTA over a window and flags HUNG on no-advance past the op's declared no-progress budget (the §11.4.201(6)-(7)/§11.4.89 discipline — `kill -0` is a necessary-not-sufficient pre-filter, a stale heartbeat/non-growing log is NEVER "progressing"). A consumer blocked on an unclosed producer write-end (§11.4.201(12)) is caught by the flat heartbeat regardless of internal state.

(D) OWNERSHIP-HANDOFF-ON-STOP. Stopping/killing an owner MUST hand its long-ops to the registry (`handoff` , re-adoptable per §11.4.147) or deliberately reap them (`reaped` , reason logged) — NEVER silently orphan or accidentally cascade-kill; long-ops launch detached (own process-group) so an owner stop does not cascade-signal them, and the stop path records each owned op's disposition. The loop done-condition is not satisfied while a `handoff` op is un-adopted.

(E) SAFE REAPING. An orphaned/hung op is reaped ONLY on PROVEN staleness — holder PID dead (`kill -0` fails) OR no-progress-budget exceeded per (C) — resolved by the process's real `/proc/<pid>/cmdline` (§11.4.196(D)/§12.12), never a bare `pgrep` /substring that could match a carrier or the reaper itself; a LIVE, ADVANCING op is NEVER reaped (§9.2); every REAPED/KEPT decision + resolved evidence is logged (§11.4.180).

(F) CONSISTENCY INVARIANT SWEEP. A fast invariant check runs before every gated transition (build/flash/tag) and on a standing cadence, asserting: no duplicate owner per purpose, no orphan, no stale lock (§11.4.180), no hung op (C), no registry-vs-reality drift (live-long-op set == registry live rows — the §11.4.196(F) check generalised). Any violation is surfaced (named, actionable) and the gated transition is REFUSED until resolved. The sweep is itself a guard (§11.4.201): golden-TRUE + golden-FALSE-with-carrier fixtures, conservative-safe REFUSE on an unresolvable signal, resolved evidence printed on every refusal.

(G) UN-COMPLETABLE-BLOCKER ESCAPE. A gate that cannot reach a terminal verdict MUST degrade, after a bounded budget, to an explicit evidence-backed decision — NEVER an indefinite silent block: §11.4.102 first (hung/(C) vs duplicate/(B) vs wrong-seam/§11.4.120 re-seat-at-activation); NOT-YET-RUNNABLE → §11.4.69 `artifact_not_yet_built` (non-blocking on an intermediate artifact, blocking only on the release candidate per §11.4.135); genuinely un-completable → §11.4.101 (autonomous reversible SKIP-with-reason + tracked §11.4.197 item, ELSE a bounded §11.4.66 operator decision) — the loop keeps progressing every other non-blocked item, it does not idle, and the escape NEVER fabricates a PASS.

Honest boundary (§11.4.6). This layer reduces orchestration mess; it does NOT make product bugs impossible, does NOT make a gate correct (§11.4.108/ §11.4.201 still bind each gate), and does NOT remove the operator from genuine blocks — it converts an indefinite silent stall into a bounded, evidence-backed, operator-visible decision. A heartbeat proves ADVANCE not CORRECTNESS; the registry proves what CLAIMS to run, (F) proves registry==reality. **Reuse-not-rewrite:** it is the multitrack per-track agent registry (§11.4.147/§11.4.116/§11.4.187) generalised from per-track workers to any long-op and bound through, per §11.4.196(F) — no parallel mechanism is invented.

Classification: universal (§11.4.17) — no hardware/vendor/project literal; consumers supply their long-op set, purpose keys, no-progress budgets, and stop-hook wiring as DATA per §11.4.35. Composes §11.4.89 / §11.4.101 / §11.4.116 / §11.4.119 / §11.4.120 / §11.4.135 / §11.4.147 / §11.4.180 / §11.4.187 / §11.4.196(D) (F) / §11.4.201(1)(6)-(12) / §11.4.207 / §11.4.69 / §12.6 / §12.12 / §9.2. Propagation gate `CM-COVENANT-114-232-PROPAGATION` (literal `11.4.232`) + recommended mechanism gates `CM-LONGOP-REGISTRY-SINGLE-SOURCE` (A) / `CM-LONGOP-SINGLE-OWNER-PER-PURPOSE` (B) / `CM-LONGOP-LIVENESS-HEARTBEAT` (C) / `CM-LONGOP-OWNERSHIP-HANDOFF-ON-STOP` (D) / `CM-LONGOP-SAFE-REAP-PROVEN-`

STALE (E) / CM-ORCHESTRATION-CONSISTENCY-SWEEP (F) / CM-UNCOMPLETABLE-BLOCKER-ESCAPE (G) + paired §1.1 mutations (run two same-purpose sweeps with no reuse/supersede → (B) FAILs; infer liveness from `kill -0` alone / drop the heartbeat-delta check → (C) FAILs; a stop that leaves an owned op with no terminal/handoff record → (D) FAILs; reap a live-advancing op OR reap by bare- `pgrep` substring → (E) FAILs; a live long-op with no registry row passes the consistency sweep → (F) FAILs; an un-completable gate blocks with no escape/decision → (G) FAILs; strip the literal → propagation FAILs; gate-code = separate work item, NOT claimed shipped §11.4.6). No escape hatch — no `--allow-duplicate-longop` , `--infer-liveness-from-pid` , `--silent-orphan-on-stop` , `--reap-live-op` , `--skip-consistency-sweep` , `--block-without-escape` , `--longop-without-registry` flag.

EXTENSION landed 2026-07-26 — §11.4.182 model-in-label + flawless alias/provider ↔ model consistency + real-time always-in-sync (User mandate, 2026-07-26). Verbatim operator mandates (2026-07-26): "Extend our logging to present models used since we will now have dynamic models choice and switching! From (example) (T1/main - claude4) into: (T1/main - claude4 - model_used), model MUST be dynamically obtained so we present name equivalents - short names - like: haiku, opus, sonnet, fable. However we MUST fully support all Claude Toolkit providers aliases, so real model used names MUST BE obtained in this case(s) as well! Examples: (T1/main - claude4 - haiku), (T1/main - claude4 - fable), (T1/main - claude4 - opus), (T1/main - claude4 - deepseek), (T1/main - claude4 - xiaomi) ... Labels MUST BE everywhere ALWAYS up to date in real time and in sync!" and "Adding model names into labels considers that providers names / aliases will be up to date too and in relation to model used since models belong to providers / aliases !!! Consistency MUST BE flawless!" This EXTENDS the §11.4.182 label discipline (it does NOT re-mint the anchor — the literal `11.4.182` and its propagation gate `CM-COVENANT-114-182-PROPAGATION` are UNCHANGED and stay GREEN; this carrier block cites §11.4.182, it is not a second §11.4.182 block-start per §11.4.227(B)). (1) **LABEL FORM EXTENDED** — the mandated §11.4.182 label is extended from `(T<N>/<branch> - <alias>)` to `(T<N>/<branch> - <alias> - <model> - <effort>)` (the trailing `<effort>` field added 2026-07-26 per §11.4.231 clause (F) — see the EFFORT-FIELD clause (7) below), `<model>` recording which model is CURRENTLY answering for that alias: for a native Claude alias, the SHORT name (`opus` / `sonnet` / `haiku` / `fable` , case-insensitive-substring-derived from the model id the agent platform reports for the most recent

turn — DYNAMIC, obtained live at label-emission time, never a fixed/assumed value since native aliases have no persistent per-alias model and the serving model can legitimately switch mid-session per §11.4.231); for a claude-toolkit PROVIDER alias (`deepseek` , `xiaomi` , `kimi-for-coding` , `opencode` , ...) the REAL provider model id as declared in that provider's host-config registry (the SAME value the alias's own launch mechanism uses — a single source of truth, never re-derived). (2) **MODEL-FIELD DERIVATION (§11.4.6, never guessed)** — a closed two-case split on the SHAPE of the already-derived `<alias>` : (a) native-alias → read the model id from the CURRENT session's own transcript (the platform's per-project per-session transcript; LAST assistant-turn's reported model → short name); (b) provider-alias → read the model id from the host multi-account-toolkit provider registry (`CMA_PROVIDER_MODEL`). Either lookup that cannot resolve yields the honest `?` — NEVER a fabricated or last-known value; an id obtained but matching none of the four native short-names is printed VERBATIM (real-but-unrecognised is more honest than a false `?`). The 3-field legacy form (no `<model>`) remains a VALID, ACCEPTED label during migration — the enforcing guard-hook `LABEL_RE` admits both forms, and NEVER treats an honest `?` model, or an absent model field, as a block-cause. (3) **FLAWLESS ALIAS/PROVIDER ↔ MODEL CONSISTENCY INVARIANT (operator NOTE, load-bearing)** — the `<alias>` / provider field and the `<model>` field MUST be MUTUALLY CONSISTENT AT ALL TIMES: the model shown MUST be a model that genuinely BELONGS to the alias/provider shown (models belong to providers/aliases), and the alias/provider name MUST itself be current — never a stale alias beside a fresh model, never a model that does not belong to the shown alias. This is guaranteed BY CONSTRUCTION, not by coincidence: the native model is read from THAT alias's own live session transcript, and the provider model is read from THAT provider's own registry entry — so the derived `<model>` is definitionally a model of the shown `<alias>` ; a derivation that ever emits a model not belonging to the shown alias is a §11.4.6 / §11.4.182 consistency violation, never a valid label. (4) **REAL-TIME, ALWAYS-IN-SYNC (§11.4.229/§11.4.106(F))** — the label is STATELESS and RE-COMPUTED at EVERY emission (every human-invoked labeler run AND every synchronous PreToolUse guard-hook firing on every Agent/Task/TaskCreate dispatch) against the LIVE environment (current `CLAUDE_CONFIG_DIR` , current session id, current transcript tail, current provider `.env`) — so a mid-session model switch (`/model` , an autonomous §11.4.231 escalation) is reflected in the VERY NEXT emission with no cache to go stale: real-time BY CONSTRUCTION, not by a sync daemon.

Honest boundary (§11.4.6): the live label has NO persisted artifact that could drift, so the Docs-Chain (§11.4.106) sync engine — which propagates PERSISTED documents/DBs — does not apply to the stateless label itself; the operator's "use/extend Docs Chain for real-time responsiveness" ask attaches to any PERSISTED label-history / model-switch LOG surface (a §11.4.197 follow-up to be tracked), which, if adopted, IS §11.4.106-Docs-Chain-synced and §11.4.208-request-history-adjacent — never faked as already-shipped. (5) **HONEST SUBAGENT-INTROSPECTION LIMIT (§11.4.6/§11.4.197)** — the native-alias transcript lookup is proven reliable when invoked from the DISPATCHING process's own shell (the guard-hook's actual execution context at dispatch time); it is NOT proven reliable when a background-dispatched subagent introspects its OWN model from inside its own subprocess on hosts where the subagent shares its parent's session-transcript with no subagent-vs-conductor marker — recorded as an open limitation to be tracked, never silently claimed solved. (6) **DIVERGENT-LABELER RECONCILIATION (§11.4.227)** — the LIVE labeler wired via `.claude/settings.json` is `constitution/scripts/multitrack/track_branch_label.sh` (both it and the top-level `scripts/multitrack/` duplicate carry the model-field extension (functionally equivalent on the alias+model fields, not byte-identical — the constitution copy additionally carries the 2026-07-26 effort field the top-level copy does not yet mirror)); the two copies are reconciled to a single source of truth (top-level a thin wrapper exec-ing the constitution copy, §11.4.177 inherited-by-reference) as a reconciliation to be tracked, never a silent divergence. (7) **EFFORT FIELD (added 2026-07-26 per the operator IMPORTANT mandate — verbatim: "Besides dynamic determining proper models to use, we MUST decide the effort as well and switch between the effort settings same as we do now for chosen model! Labeling MUST BE extended for this details as well. Example: (T1/main - claude4 - opus - high).")**. The 5-field label appends `<effort>` after `<model>`, recording which reasoning-EFFORT setting is CURRENTLY in force for that work-stream, drawn from the closed set `{low | medium | high | xhigh | max}` (the SAME effort ladder §11.4.231 clause (F) selects). (a) **DERIVATION (§11.4.6, never guessed)** — `<effort>` is read from whatever LIVE effort signal the execution path exposes (the Workflow-tool `agent()` `effort` argument, a future Agent-tool effort parameter, an env/config effort setting); a path exposing NO effort signal, OR any value outside the closed set, yields the honest `?` — EXACTLY like the `<model>` field's `?` fallback, NEVER a fabricated / assumed /

last-known value. (b) HONEST HARNESS BOUNDARY (§11.4.6, load-bearing — do NOT overstate the mechanism) — effort-tiering is REAL where the execution path exposes it and a DOCUMENTED CAPABILITY-GAP where it does not: the current Claude Code **Agent tool** (subagent dispatch) exposes a `model` parameter but **NO** `effort` parameter, so a subagent's effort is NOT currently settable via that path and its `<effort>` field degrades honestly to `?`; the **Workflow tool's** `agent()` DOES take an `effort` argument, so effort IS settable + derivable there — the label reports whichever path is live, never invents one. This is the SAME aspirational-vs-harness state the §11.4.209 / §11.4.211 Fable- `xhigh` review/merge pins already occupy: the pins mandate the effort, and whether a given dispatch path can SET it is a harness capability the label reports honestly and NEVER fakes. (c) LEGACY-FORM ACCEPTANCE (migration) — the 3-field (`T<N>/<branch> - <alias>`) and 4-field (`T<N>/<branch> - <alias> - <model>`) label forms remain VALID and ACCEPTED; the guard `LABEL_RE` admits all three shapes and NEVER treats an absent `<effort>` field, or an honest `?` effort, as a block-cause (§11.4.6 — mirroring the §11.4.182 track/alias/model honest- `?` rule). (d) CONSISTENCY — the effort field is selected TOGETHER WITH the model by §11.4.231 (both tiers chosen for one task) and never contradicts it; a stale effort beside a fresh model, or an effort outside the closed set, is a §11.4.6 / §11.4.182 consistency violation, never a valid label. (e) REAL-TIME, ALWAYS-IN-SYNC (§11.4.229) — like `<model>`, `<effort>` is STATELESS and RE-COMPUTED at EVERY label emission against the LIVE environment, so a mid-task §11.4.231 effort escalation is reflected in the VERY NEXT emission with no cache to go stale. Propagation: `CM-TRACK-BRANCH-LABEL` is EXTENDED to also assert the labeler emits a parseable `<model>` field AND a parseable `<effort>` field, the guard `LABEL_RE` accepts the 3-field / 4-field / 5-field label forms (legacy 3-field and 4-field ALWAYS remain accepted during migration; a `?` effort is NEVER a block-cause), and the alias/provider ↔ model consistency holds; paired §1.1 mutation: strip the `<model>` derivation OR emit a model not belonging to the shown alias → the extended gate FAILs (gate-code = separate work item, NOT claimed shipped §11.4.6). No new escape hatch — the existing §11.4.182 "no escape hatch" list applies to the extended label unchanged.

§11.4.233 — Standing anti-mess control plane: a level-triggered, idempotent, invariant-catalogue reconciliation loop over the WHOLE persistent System state, orchestrating every per-domain enforcer, refusing gated transitions on un-reconciled drift (research-derived, 2026-07-26).

Forensic anchor (genericised): every existing consistency enforcer in a multi-track / multi-agent / multi-submodule / multi-upstream System is EDGE-TRIGGERED — a PreToolUse hook fires on an action, a check runs at a seam (lock-acquisition / merge-time / build-time / commit-time), a gate runs when a build runs — so a mess that arises WITHOUT an edge firing is un-observed until it causes harm (two live examples captured 2026-07-26: a large uncommitted-batch pile-up sat unreconciled through a session, and multiple provably-stale VCS/wrapper locks were held by dead holders with no acquisition attempt to reap them — the §11.4.205 "a hook never installed / a check that runs only when asked" forensic generalised to the whole persistent state). Per §11.4.102 Phase 4.5 the ~19 persistent-state mess classes (Plane R repository: uncommitted pile-up, stale/duplicate locks, orphaned worktrees, half-applied cross-repo commits, drifted submodule pointers, partial merges, divergent mirrors, mis-placed work; Plane G governance: divergent mirrors, pointer drift, gate-ledger drift, anchor-number collision; Plane S tracking: doc/DB drift, cross-track SSoT divergence, ledger drift, handoff-doc drift; Plane P runtime: the §11.4.232 long-op set + un-reaped streams + dangling review state) are ONE architectural flaw — there is no LEVEL-TRIGGERED control plane that re-derives the actual persistent state and reconciles it against a declared clean state; an un-observed drift is a §11.4/§11.4.1 bluff at the organization layer ("the System is clean" while a mess silently accumulates).

The mandate — a standing **anti-mess control plane** (the desired-state/actual-state reconciliation-loop + control-plane/data-plane pattern of distributed systems), ALL of which hold:

(A) DECLARED INVARIANT CATALOGUE. A single machine-readable consumer-owned catalogue (DATA per §11.4.28/§11.4.35) enumerating every consistency invariant across the persistent planes, each carrying `{id, plane, desired-condition, detector, reconcile-class ∈ {auto-safe | operator-gated | owned-by <anchor>}, evidence-shape}` ; the catalogue IS the desired clean state; an un-catalogued mess class is an honest tracked GAP (§11.4.6/§11.4.197), never silently claimed covered.

(B) LEVEL-TRIGGERED LOOP. The loop re-derives the ACTUAL state (not an event) and diffs it against the catalogue BEFORE every gated transition (commit-batch/build/flash/merge-to-trunk/tag/constitution-pull) AND on a standing cadence

(§11.4.94/§11.4.103) — catching drift regardless of how it arose (the Kubernetes level-triggered property); actual-vs-desired uses a §11.4.86 fingerprint diff drilling to the offending record (anti-entropy), never a whole-scan where a fingerprint suffices.

(C) IDEMPOTENT RECONCILE-OR-REFUSE. Each drift takes exactly one dedup-keyed action: RECONCILE-AUTO-SAFE (reversible + evidence-determinable + bounded-blast per §11.4.101 — reap a provably-dead lock §11.4.180, regen a stale derived doc §11.4.106, scope-group-commit a placement-clean batch §11.4.191, retire a merged worktree §11.4.167), idempotent (a no-op on a clean state); OR REFUSE-THE-GATE + a bounded operator decision (irreversible/high-blast/indeterminate — id-colliding SSoT §11.4.206, anchor collision §11.4.227, partial cross-repo merge, divergent mirror) per §11.4.101/§11.4.66; a gated transition NEVER proceeds while a catalogued invariant is violated and un-reconciled; the loop keeps progressing every other non-blocked item (park-one-not-the-loop).

(D) SAGA FORWARD-RECONCILE (no-loss). A partially-applied cross-repo operation is reconciled complete-forward under a §9.2 pre-op backup (finish the remaining legs / bump the pointer / complete the merge zero-markers-zero-drop §11.4.41 step 3 / ff the lagging mirror) — NEVER force-push/reset/ -s ours (§11.4.113/§9.2); compensation is §9.2 backup-restore only; the saga reaches fully-completed, never rolled-back-with-loss.

(E) THE LOOP IS A GUARD (§11.4.201). Every detector ships golden-TRUE + golden-FALSE-with-carrier fixtures (§11.4.107(10)); a false-positive REFUSAL of a clean state is a §11.4.201(1) FAIL-bluff exactly as a false-negative PASS is a §11.4 PASS-bluff; every "no drift" null is control-needle-proven (§11.4.201(6)-(12)); conservative-safe REFUSE on an unresolvable signal, resolved evidence printed on every refusal.

(F) CONTROL-PLANE / DATA-PLANE SEPARATION. The control plane DETECTS + RECONCILES state and NEVER performs the data-plane work (the actual build/commit/merge/long-op — §11.4.232 + the build/flash/commit wrappers); the two never merge; the control plane's own state is durable (§11.4.205(6) temp → fsync → rename, never tmpfs, may reuse §11.4.207).

Honest boundary (§11.4.6). Converges the CATALOGUED invariants; does NOT make an un-catalogued mess class impossible (extensible catalogue, never "provably mess-free"), does NOT make a gate correct (§11.4.108/§11.4.201), does NOT auto-fix irreversible drift (operator-gated §11.4.101/§11.4.66), and is eventual-

consistency-at-every-gate-and-cadence NOT a real-time daemon (§11.4.205(7)).

Reuse-not-rewrite: it is the level-triggered COMPOSER over the existing per-domain enforcers (§11.4.180/§11.4.186/§11.4.106/§11.4.191/§11.4.206/§11.4.227/§11.4.232/§11.4.147/§11.4.167), inheriting each's detector — no parallel enforcer is invented (§11.4.227). It composes §11.4.232 (the RUNTIME data-plane long-op orchestrator — §11.4.232 clause (F) is Plane P of this catalogue) as the plane it does not touch.

Classification: universal (§11.4.17) — no hardware/vendor/project literal; consumers supply their invariant catalogue, plane detectors, gated-transition list, cadence, and reconcile-class bindings as DATA per §11.4.35. Composes §11.4.6 / §11.4.26 / §11.4.28 / §11.4.32 / §11.4.84 / §11.4.86 / §11.4.101 / §11.4.106 / §11.4.108 / §11.4.116 / §11.4.121 / §11.4.147 / §11.4.164 / §11.4.167 / §11.4.176 / §11.4.179 / §11.4.180 / §11.4.181 / §11.4.186 / §11.4.191 / §11.4.196(F) / §11.4.201 / §11.4.205 / §11.4.206 / §11.4.207 / §11.4.227 / §11.4.232 / §12.6 / §9.2 / §2.1. Propagation gate CM-COVENANT-114-233-PROPAGATION (literal 11.4.233) + recommended mechanism gates CM-ANTIMESS-INVARIANT-CATALOGUE (A) / CM-ANTIMESS-LEVEL-TRIGGERED-LOOP (B) / CM-ANTIMESS-RECONCILE-OR-REFUSE (C) / CM-ANTIMESS-SAGA-FORWARD-NO-LOSS (D) / CM-ANTIMESS-LOOP-IS-A-GUARD (E) / CM-ANTIMESS-CONTROL-DATA-PLANE-SPLIT (F) + paired §1.1 mutations (drop the standing-cadence trigger so the loop is edge-only → (B) FAILs on the pile-up/stale-lock fixture; let a gated transition proceed on a catalogued un-reconciled drift → (C) FAILs; a saga that force-pushes/resets to reconcile a mirror → (D) FAILs; a detector that REFUSES a clean golden-TRUE fixture, or PASSES a golden-FALSE-with-carrier → (E) FAILs; the control plane invokes a data-plane build/commit → (F) FAILs; strip the literal → propagation FAILs; gate-code = separate work item, NOT claimed shipped §11.4.6/§11.4.227). No escape hatch — no --edge-triggered-only, --skip-consistency-catalogue, --proceed-on-drift, --reconcile-with-force-push, --unvalidated-detector, --control-plane-does-work, --assume-clean-without-reconcile flag.