

QWEN

**Helix Constitution – Universal
QWEN.md**

Field	Value
Revision	53
Created	2026-05-20
Last modified	2026-07-17T14:30:00Z
Status	active

Status summary

Round 2 of 2026-07-17 — the CARRIER / MEASUREMENT-INTEGRITY harvest (research-derived from one cycle's forensics in a consuming project). 4 EXTENSIONS + 2 NEW anchors + 1 candidate REJECTED-as-already-covered. The unifying FACT: every false finding that cycle had ONE shape — a token that MENTIONS X matched as X, or an instrument's silence read as data (a removal's ref-count matched the audit-trail comments DOCUMENTING the removal; escapes eaten by a wrapper shell so the far side searched a literal; `\|` a literal pipe under ERE; `check |` head returning the PAGER's exit status so a parse-check said 'clean' 3x without checking; a line-anchored pattern missing an indented colour-escaped verdict; a binary extractor splitting a multi-byte char) — and §11.4.196(D)/§12.12 had ALREADY anchored the shape TWICE while it kept recurring; the self-demonstrating instance: a mutation-residue scan flagged the very anchor that DEFINES the markers because it quotes them (proven carrier by identical HEAD-vs-worktree counts). The in-session fix: the CONTROL NEEDLE. §11.4.201(6)(7)(8) — scope widened from GUARDS to EVERY load-bearing measurement/metric with a MEASUREMENT-shaped taxonomy (false-null / false-match; the null more dangerous — a blind instrument and a clean artifact return the identical quiet zero); match STRUCTURE not substring; a NULL is not evidence until a

control needle proves the instrument can see through the SAME path (generalises §11.4.115(F) validate-the-detector, PER-QUERY — goldens necessary-not-sufficient §11.4.146(D3)); THE PATH IS PART OF THE INSTRUMENT (quoting/dialect/pipeline-exit/anchoring/encoding all return clean confident wrong answers without crashing, so §11.4.1 and §11.4.67 cannot see them); and a METRIC MUST be validated against the DEFINITION OF DONE — if the correct end-state moves it the wrong way it is INVALID and must not gate work (honest boundary: §11.4.135's adoption ratchet is one such metric, its validation OWED not claimed, tracked §11.4.197).

§11.4.112(5) — impossibility verdicts are BOUNDED: state the EXACT scope, ENUMERATE adjacent goals NOT covered, and be CITATION-FENCED (a TRUE verdict leaked to an adjacent VIABLE goal for six weeks while a public SDK call the vendor's own docs prescribed existed; the anchor was asymmetric — clause (3) hardened against reopening, nothing bounded REACH; §11.4.150 fires where a verdict is MINTED, this leak happened where one was CITED). §11.4.120(seam-placement) — a gate whose precondition the gated work itself produces is on the WRONG SEAM (belongs at activation/release, never authorship); re-read the source of truth before re-seating; the report-don't-improvise half is already §11.4.6/.66/.101/.201(4). §11.4.69 — `artifact_not_yet_built` added to the closed reason-set for the §11.4.135

REQUIRES-REBUILD PENDING class (seam-dependent per §11.4.135, NOT flat). REJECTED as already-covered: 'BLOCKED as a third verdict state' — §11.4.135's verdict-coverage extension states it MORE precisely and explicitly rejects the flat framing the candidate proposed; only the reason-code was genuinely missing. NEW §11.4.214 — recurrence-links-not-mints: a returning defect MUST reopen its existing item, never mint a new id (recurrences as new ids silence the §11.4.55 reopens-count — the exact signal §11.4.132(d)/ §11.4.189 use to rank the most-fragile work — precisely on the items that keep breaking; the recorded 52% rate was a known UNDERCOUNT; ~15 items mapped to ~9 root causes); every intake path bound, §11.4.186 keying + negative-control, distinct-but-similar a FIRST-CLASS verdict, autonomous default mint-with-link over merge (a spurious id is recoverable, a wrongly-merged defect is LOST). NEW §11.4.215 — a doc that BINDS work MUST live TRACKED in the repo where that work happens (the doc DEFINING a binding gate was untracked and out-of-tree); binding⇒tracked, un-tracked⇒NOT binding; the PRIOR question §11.4.212 cannot ask (it enumerates the in-repo scope, so a doc in no checkout is never enumerated); the generalisation of §11.4.95 exactly as §11.4.212 generalised §11.4.57; §11.4.11 intact — binding is what pulls a doc into the repo. All Classification: universal (§11.4.17). Propagation gates CM-

COVENANT-114-214/215-
PROPAGATION; literals
11.4.201/.112/.120/.69 unchanged (gates
stay GREEN). CLAUDE.md +
AGENTS.md + QWEN.md + GEMINI.md
mirrors updated in lockstep (§11.4.157).
Prior round: Strengthened §11.4.115 +
§11.4.135 + §11.4.146 (research-
derived, 2026-07-17 — the status-
without-custody forensics from a
consuming project): NO new anchor
minted — the rules already existed as
prose and did not bind; landed the
MECHANICAL BINDING instead.
Forensic FACTs (genericised): 576
tracked items, 182 done-claiming, 173
(95%) with NO registered guard — an
operator sampled 6 done-claimers, 6/6
broken (the expected draw of that
arithmetic, not bad luck); recorded
reopen-per-recorded-fix 52% (elite
<10%), itself an UNDERCOUNT
because recurrences re-entered as new
ids and statuses moved with zero history
rows; the guard suite's exit code read
ONLY its FAIL count — measured live on
a release candidate: 61 guards → PASS
0 / FAIL 5 / PENDING 56 → exit 0 “not
blocked” — and the release-tag tooling
consulted the guards ZERO times, so
“never verified” and “verified good” were
the same colour; the project's own
history selected the discriminator: every
fix whose done-claim was preceded by
an on-device RED → GREEN flip that
HELD recorded zero reopens, every fix
confirmed at source-green bounced.
§11.4.115(F) — machine-written verdict
pairs + guard-viability: RED/GREEN are

harness-written verdict files (target fingerprint read from the target at run time, RED fingerprint $\neq$ GREEN fingerprint, iterations ≥ 3 , evidence class matches the defect's layer — a grep transcript can never satisfy a user-visible class); a guard never observed FAILing on the genuinely-broken artifact is unvalidated instrumentation and mints no verdicts; the canonical §1.1 mutation is the fix-commit's revert, and a string-deletion mutation that only removes what the gate greps is a refused tautology; gate CM-GUARD-VERDICT-MACHINE-WRITTEN . §11.4.135 verdict-coverage extension — the suite exit MUST incorporate COVERAGE, never only N_FAIL; the release seam blocks on `uncovered = registered ^ topology-present ^ no-verdict-for-the-candidate-fingerprint` (ABSENCE of a verdict blocks exactly as a FAIL does); the release-tag tooling MUST consult the verdict store; REQUIRES-REBUILD PENDING is self-clearing on the candidate, topology exemptions enumerated via a checked-in map never silent, adoption via a monotone-decrease ratchet, and a naive “PENDING always blocks” is rejected as a §11.4.201 FAIL-bluff; gate CM-RELEASE-VERDICT-COVERAGE . §11.4.146(D3) — status-custody: a done/ready status write is REFUSED at the status-write seam without the file chain registry-row-keyed-by-the-item's-EXACT-id (alias map for legacy keys) → executable guard → RED+GREEN verdict pair → class-matched evidence

(prose nowhere in the chain); a full-table sweep gate (not a diff check) catches raw-DB bypasses; §11.4.205 construction order — the executable sweep + its golden-bad fixture land in the SAME commit as any governance text citing them, and the meta-test asserts the sweep is WIRED; gate CM-STATUS-CUSTODY . Honest boundary §11.4.6 carried explicitly: the binding does not prove features bug-free (§11.4.186 family dedup remains), does not fix a miscalibrated oracle (§11.4.107(10) goldens necessary-not-sufficient), does not replace §11.4.185 manual QA, and makes target availability a hard dependency BY DESIGN. All Classification: universal (§11.4.17) — no hardware/vendor/project literal; consumers supply registry/sweep/topology-map paths as DATA per §11.4.35. Propagation gates for literals 11.4.115 / 11.4.135 / 11.4.146 unchanged (stay GREEN). CLAUDE.md + AGENTS.md + QWEN.md + GEMINI.md mirrors updated in lockstep (§11.4.157). Prior round: Added §11.4.213 — FEATURE research-scheduling directive (operator-directed mandate, 2026-07-16): a new §11.4.140-registered action FEATURE (all six grammar forms incl. the single-colon FEATURE: form) that SCHEDULES — never synchronously runs — a deep, enterprise-grade research + implementation-planning effort as a Type=Task workable item, reusing the §11.4.202 report_item.sh engine for item-creation + DB-sync + tracker-push

(never reimplemented) plus a new durable docs/requests/feature_queue.md queue so a scheduled request is never dropped; the created item embeds the full mandate (deep web research §11.4.8/.99/.150; systematic-debug of every weak-spot/gap/danger-zone §11.4.102 with a risk-free rock-solid design per gap; always enterprise-grade/bleeding-edge/innovative; exhaustive documentation down to lines-of-code + micro-POCs + diagrams + SQL + templates §11.4.65/.73; reuse-first investigation of vasic-digital/HelixDevelopment components before rewriting, kept decoupled §11.4.28/.74/.177; full test-type + Challenges + HelixQA planning §11.4.27/.169; fine-grained phased/task/subtask breakdown; enterprise-scalability + max-performance planning; OpenDesign UI/UX wireframes where applicable §11.4.162/.190; full CodeGraph integration §11.4.78/.79/.80; fully-detailed workable items synced in real time to the SQLite SSoT + every connected external tracker §11.4.93/.95/.148 D5, honest-skipping any absent tracker §11.4.10) as its acceptance-defining work program — the multi-day research itself is EXECUTED by the standing autonomous loop (§11.4.87/.94/.97/.103/.126) when it claims the item, driven to a genuinely COMPLETED-and-wired or explicitly evidence-backed CLOSED terminal state under the §11.4.197 research-completion mandate, NEVER left sitting un-wired in

Field	Value
	the backlog. Propagation gate CM-COVENANT-114-213-PROPAGATION (literal 11.4.213) + recommended gate CM-FEATURE-DIRECTIVE + paired §1.1 mutation. Classification: universal (§11.4.17). CLAUDE.md + AGENTS.md + QWEN.md + GEMINI.md mirrors updated in lockstep (§11.4.157). Prior round: Added §11.4.211 + §11.4.212 (User mandate 2026-07-16). §11.4.211 — Merge-conflict resolution during any main → feature/product/flavor merge MUST run on the Fable model at xhigh effort (Opus xhigh fallback): the merge-CONFLICT-resolution analogue of §11.4.209's code-REVIEW model-pin, generalised per §11.4.17 to EVERY track's controlled main → (feature
Issues	none
Issues summary	—
Fixed	§11.4.108 mirror
Fixed summary	§11.4.107 lands in lockstep with the Constitution.md §11.4.107 addition.
Continuation	—

Table of contents

- [How inheritance works](#)
- [Identity & posture](#)
- [Top-level invariants for every agent](#)

- Critical base rules restated (for agents that don't honour @imports)
 - §11.4 — Anti-bluff covenant — END-USER QUALITY GUARANTEE
 - §1.1 — Mutation-paired gates
 - §11.4.10 — Credentials-handling mandate
 - §11.4.17 — Universal-vs-project classification
 - §11.4.20 — Subagent-driven-by-default
 - §11.4.73 — Main-specification document versioning + revision discipline
 - §11.4.74 — Submodule-catalogue-first discovery + extend-don't-reimplement
 - §11.4.76 — Containers-submodule mandate
- Companion documents

How inheritance works

This `QWEN.md` is for the **Qwen Code CLI agent** (`qwen-code`). It documents the universal Helix Constitution from the Qwen agent's perspective — the same content seen by Claude Code via `CLAUDE.md` and by OpenCode / Cursor / generic AI tooling via `AGENTS.md` .

The **authoritative source** is `Constitution.md` in this repository. When this file diverges from `Constitution.md` , the latter wins. This file exists because:

1. Qwen Code does not always honor `@import` directives across files; restating the critical rules inline ensures the agent reads them on every session.
2. Cross-agent parity: Claude Code (`CLAUDE.md`), OpenCode/Cursor (`AGENTS.md`), Qwen Code (`QWEN.md`) all start from the same constitutional baseline.

Discover this file via the canonical parent-walk pattern documented in every consuming project's `CLAUDE.md` / `AGENTS.md` :

```

find_constitution_root() {
  local cur="${1:-$PWD}"
  while [[ "${cur}" != "/" ]]; do
    if [[ -f "${cur}/constitution/Constitution.md" ]]; then
      echo "${cur}/constitution"; return 0
    fi
    cur="$(dirname "${cur}")"
  done
  return 1
}

```

Identity & posture

When the Qwen Code agent loads this file as part of session bootstrap, it operates under the Helix Constitution with the same identity, anti-bluff posture, and end-user quality covenant as any other CLI agent in the catalogue. There is no “Qwen-lite” interpretation — Qwen Code is held to the full §11.4 covenant.

Top-level invariants for every agent

1. **Read order on a cold start.** `CLAUDE.md` / `AGENTS.md` / `QWEN.md` (this file) for the AI-agent posture; `Constitution.md` for the canonical universal rules; then the consuming project’s `CLAUDE.md` for project-specific extensions.
2. **Multi-mirror push.** Every commit on this submodule MUST land on all four canonical mirrors (GitHub + GitLab + GitFlic + GitVerse) per §2.1.
3. **Anti-bluff is non-negotiable.** Passing tests MUST imply working features. Bluffing PASSes (and FAILs that hide root cause) is a release blocker (§11.4 + §11.4.1).
4. **Submodule-catalogue-first.** Before scaffolding anything, survey the `vasic-digital` + `HelixDevelopment` orgs for an existing Submodule (§11.4.74). Reuse → extend → no-match in that order.
5. **Containers-submodule for ALL container workloads.** Use `vasic-digital/containers` (§11.4.76) — never reinvent compose/boot orchestration in-project.

Critical base rules restated (for agents that don't honour @imports)

§ . – Anti-bluff covenant – END-USER QUALITY GUARANTEE

Forensic anchor — verbatim user mandate (2026-04-28, reasserted 2026-05-21, 2026-05-29):

“all existing tests and Challenges do work in anti-bluff manner - they MUST confirm that all tested codebase really works as expected! We had been in position that all tests do execute with success and all Challenges as well, but in reality the most of the features does not work and can't be used! This MUST NOT be the case and execution of tests and Challenges MUST guarantee the quality, the completion and full usability by end users of the product!”

Operative rule. Captured tests MUST exercise the behavior they claim to verify. PASS-without-execution paths, mock-only tests that don't round-trip, skip-by-default integration tests that mask real failures, metadata-only PASS, configuration-only PASS, “absence-of-error” PASS, and grep-without-runtime PASS are all critical defects regardless of how green the summary line looks. Each test failure in CI MUST imply a real broken feature; each PASS MUST imply the feature actually works for the end user. Tests and HelixQA Challenges are bound EQUALLY.

The verbatim covenant MUST be present in every consumer governance file (project Constitution.md, CLAUDE.md, AGENTS.md, QWEN.md) and in every owned submodule's equivalent. Tools that don't expand `@imports` still read the literal text.

§ . – Mutation-paired gates

Every captured-evidence gate MUST ship with a paired mutation test that mutates the gate's anchor + runs the gate + asserts the gate now FAILs. Absence of the paired mutation is itself a bluff.

§ . . – Credentials-handling mandate

Credentials are NEVER committed to git, NEVER logged, NEVER printed to stdout/stderr. `.env` files are `.gitignore` d; committed siblings are `.env.example` placeholders only. Pre-store leak audit (§11.4.10.A) scans every PR for credential patterns before merge.

§ . . – Universal-vs-project classification

Every new rule MUST declare itself **universal** (applies to every project consuming this Constitution) or **project** (applies only to one specific project). Misclassification (universal rule landing in a project's local CLAUDE.md instead of the constitution submodule) is a release blocker.

§ . . – Subagent-driven-by-default

When a CLI agent (Qwen Code included) has a subagent / Task tool available, the default mode of operation is to dispatch subagents for discrete units of work rather than perform them inline. This protects context, enables parallel work, and matches the §11.4.27 no-fakes-beyond-unit-tests posture (each subagent runs real tools).

§ . . – Main-specification document versioning + revision discipline

Projects with a main specification (`docs/specs/.../specification.V<N>.md` -shaped) maintain TWO version axes:

- **Primary** (V1 / V2 / V3 ...) bumps for major rewrites only.
- **Secondary** (`Revision N` in the metadata table) bumps for additive changes within a primary version.

Spec edits trigger mandatory comprehensive planning and implementation in the consuming project — they are not isolated doc tweaks (the “spec ripple” rule).

§ . . – Submodule-catalogue-first discovery + extend-don't-reimplement

Direct user mandate (verbatim, 2026-05-20): “We MUST ALWAYS check which already developed features / functionalities do exist as a part of our comprehensive Submodules catalogue located in `vasic-digital` and `HelixDevelopment` organizations on GitHub and GitLab both! Project MUST BE aware of all its existence so we do not implement same things multiple times. For any missing features we MUST IMPLEMENT them properly and extend those Submodules further!”

Before scaffolding any new module / package / helper / utility, the Qwen Code agent MUST: (1) survey `vasic-digital` + `HelixDevelopment` on GitHub + GitLab — the canonical inventory is `submodules-catalogue.md` (142 repos categorised); (2) reuse an existing Submodule when it covers $\geq 80\%$; (3) extend in-place via upstream PR when 80%+ matches; (4) document the survey result via `Catalogue-Check: reuse|extend|no-match <org/repo>@<sha>` in the tracker row.

Every Submodule in the catalogue is subject to the same development-cycle rules (§11.4 anti-bluff, §1.1 paired mutations, §11.4.10 credentials, §11.4.61 metadata + ToC, §11.4.65 universal export, §11.4.73 spec versioning, §2.1 multi-mirror push, §3 propagation order).

§ . . – Containers-submodule mandate

Direct user mandate (verbatim, 2026-05-20): “For any work or requirements of running services or codebase inside the Containers (Docker / Podman / Qemu / Emulators, and so on) we MUST USE / INCORPORATE the Containers Submodule properly: `https://github.com/vasic-digital/containers` (`git@github.com:vasic-digital/containers.git`). Containers Submodule contains all means for us to Containerize our code and services! If any feature or Containing System is missing or not supported we MUST EXTEND IT properly like we do all of our projects! No bluff work is allowed of any kind!”

For ANY containerized workload (Docker / Podman / Qemu / Kubernetes / container-backed emulators), the Qwen Code agent MUST: (1) install `vasic-digital/containers` (`digital.vasic.containers`) as a Git submodule when scaffolding the consuming project; (2) consume via `replace` directive during

development + pinned commit SHAs in production; (3) boot infra on-demand via the Submodule's `pkg/boot` + `pkg/compose` + `pkg/health` APIs so the operator is never required to start `podman machine` / `docker compose up` manually — the boot is part of the test entry point (**on-demand-infra invariant**); (4) extend the Submodule via upstream PR when a runtime / lifecycle primitive is missing — never reimplement in-project (per §11.4.74); (5) anti-bluff: integration tests claiming to exercise containerized components MUST actually boot them via the Submodule. Short-circuit fakes that bypass boot are a §11.4 violation. A passing test MUST imply the infra was up.

Tracker rows touching containerization MUST record `Catalogue-Check: extend vasic-digital/containers@<sha> (or reuse )`; `no-match` requires demonstrating the Submodule cannot model the workload.

Planned anti-bluff gate `CM-CONTAINERS-USED` scans container-touching PRs for `digital.vasic.containers/...` imports. Paired mutation strips the import + asserts FAIL.

Canonical authority: `Constitution.md` §11.4.76.

Non-compliance: reinventing compose orchestration in-project is a release blocker.

§ . . — Regeneration-mechanism-required mandate

Direct user mandate (verbatim, 2026-05-20): “We must be sure that after excluding anything from Git versioning we still have the mechanism which will out of the box obtain or re-generate missing content! Add this mandatory safety rule / constraint into ours root (constitution Submodule) `Constitution.md`, `CLAUDE.md`, `AGENTS.md`, `QWEN.md` or any other required document / file from the constitution Submodule!”

Forensic anchor. 2026-05-20T15:00Z: a consuming project's audio Tier 1

`commit_all.sh` stalled 4 h on `git add -A` scanning 274 GiB `.git-backup-*` + 159 GiB `RKTools/linux/` + 167 GiB `qa-results/` — all untracked but un-gitignored. Bare `.gitignore` fix would orphan every fresh clone (missing `RKTools`, missing test infra, missing build outputs).

Every `.gitignore` entry the Qwen Code agent considers adding that excludes (a) >~100 MiB OR (b) any artefact essential to building / running / testing the project MUST carry a documented + automated mechanism to either **re-obtain** (vendor tarball, SDK installer, package registry, git submodule, object store) OR **re-generate** (build pipeline, code-gen, asset render, captured-evidence replay, container build).

Required artefacts per qualifying `.gitignore` entry: (1) `.gitignore-meta/<entry-slug>.yaml` declaring pattern + mechanism-type + script-path + expected-disk-usage + vendor-url-or-source + integrity hash + requires-network + requires-credentials; (2) post-clone bootstrap entry (`scripts/setup.sh` or canonical equivalent) running the mechanism non-interactively; (3) pre-build gate verifying regenerated content present OR stamp `.gitignore-meta/.regenerated/<slug>.ok` recent; (4) README + `docs/guides/*.md` describing the mechanism + manual fallback + time/disk budget + per-§11.4.10 credentials.

No escape hatch: bare `.gitignore` additions without the mechanism are §11.4 PASS-bluff variants — codebase appears complete but fresh clone cannot build / run. No `--skip-regen-mechanism` , `--gitignore-is-enough` , `--operator-already-has-content` flag.

Planned anti-bluff gate `CM-GITIGNORE-REGEN-MECHANISM` scans every `.gitignore` addition for matching `.gitignore-meta/` sibling. Paired §1.1 mutation strips a required YAML key → gate FAILs.

Composes with §11.4.6, §11.4.65, §11.4.66, §11.4.71, §11.4.74, §11.4.75, §11.4.76, §9 / §9.2, §3.

Canonical authority: `Constitution.md` §11.4.77.

Non-compliance is a release blocker regardless of context.

§ . . – CodeGraph code-intelligence mandate

Direct user mandate (verbatim, 2026-05-20): “Make codegraph MANDATORY CHOICE for this purpose for all of our project ... All project which do not have configured and installed codegraph yet MUST DO IT and MUST USE IT!”

Every consuming project worked on by AI coding agents — Qwen Code included — MUST install, initialize, and use **CodeGraph** (<https://github.com/colbymchenry/codegraph> , `npm @colbymchenry/codegraph`): a local SQLite semantic code-knowledge-graph exposed to agents over MCP, 100% local. Install globally via `npm` (no `sudo`). `codegraph init` + `codegraph index` : `.codegraph/config.json` is tracked, `.codegraph/codegraph.db` is gitignored with `codegraph index` as its §11.4.77 regeneration mechanism; the `exclude` list MUST exclude every §11.4.10 credential/secret path. **CRITICAL — CodeGraph filter mechanism:** CodeGraph v1.x is zero-config — it uses built-in skip lists + root `.gitignore` for file filtering, NOT `config.json` include/exclude. Use root-anchored `.gitignore` patterns to exclude only RUBBISH (build outputs, caches, credentials). **ALL SOURCE CODE MUST BE INDEXABLE** — AOSP platform trees, third-party vendored trees, and all tracked source code are part of the index. See `Constitution.md` §11.4.78 for the complete mechanism description. Wire the `codegraph serve --mcp` server into every CLI agent the developers use — for Qwen Code via `.qwen/settings.json` (`qwen mcp add codegraph codegraph serve --mcp --scope project --transport stdio`). Cover the integration with an anti-bluff verification suite using an unforgeable per-agent challenge (a fact obtainable only by calling a CodeGraph MCP tool); un-runnable agents are documented SKIP gaps, never faked PASSes. Document in `docs/CODEGRAPH.md` . CodeGraph is consumed as the `npm` package (§11.4.74), not a git submodule.

Universal mandatory-exclude baseline set (§11.4.78 + §11.4.79 refinement — binds EVERY consuming project). Independently of project, the CodeGraph index MUST NEVER include the following file-classes anywhere — they are indexed-rubbish in every project; a consuming project MAY add its own project-specific rubbish paths on top per §11.4.35 but MUST NOT drop any baseline class: (a) **build outputs / regenerable artifacts** — `/out/` , `/build/` , `/dist/` , `/target/` , `bin/` , `obj/` , `.compressed/` , and any generator-produced tree; (b) **caches / dependency / tooling dirs** — `**/ccache/` , `**/.ccache/` , `.gradle/` , `__pycache__/` , `.pytest_cache/` , `.mypy_cache/` , `.ruff_cache/` , `node_modules/` , `.next/` , `.cache/` , `.terraform/` , `.venv/` , `logs/` ; (c) **credentials / secrets** — `.env` , `.env.*` , `**/*.env` , `secrets/` , `**/*secret*` , `keystores` / `signing-keys` / `service-account JSON`, and every credential-bearing path per §11.4.10 (a secret reaching the index is a §11.4.10

violation, never negotiable); (d) **QA / recording / large regenerable corpora** — `**/qa-results/` , `**/recordings/` , and prebuilt-binary trees. Realized via root-anchored `.gitignore` patterns (CodeGraph v1.x zero-config filter = built-in skip list + root `.gitignore` , NOT `config.json` include/exclude). Own-org submodule sources stay INCLUDED per §11.4.79; third-party vendored trees EXCLUDED. This baseline is UNIVERSAL (§11.4.17) — dropping any class is a §11.4.78 / §11.4.79 violation, and no source-code tree may be excluded to satisfy it.

Composes with §11.4.3, §11.4.10, §11.4.12, §11.4.65, §11.4.30, §11.4.74, §11.4.77, §11.4, §1.1.

Canonical authority: `Constitution.md` §11.4.78.

^{11 4 79} Non-compliance is a process violation; a project worked on by AI agents without CodeGraph installed, wired, and anti-bluff-verified is in breach.

§ . . — Own-org submodules MUST be included in the CodeGraph index

Direct user mandate (verbatim, 2026-05-21): “All Submodules we use in the project and that are part of organizations to which we have the full access via GitHub, GitLab and other CLIs MUST BE included into the codegraph database and initialized / scanned / synced!”

Refines §11.4.78 step 2 with a per-submodule-ownership split. **Own-org submodules** — full write access via the project’s CLIs (canonical orgs: `vasic-digital` + `HelixDevelopment`) — MUST be INCLUDED in the index so Qwen Code (and every other AI agent) sees a unified call-graph across the project’s domain + own-org infra code. **Third-party submodules** (the §11.4.74 `no-match` → `vendor` path) MUST be EXCLUDED. Operational steps: (1) `git submodule update --remote --merge` to pull latest before re-indexing — respect load-bearing third-party pins. (2) Adjust `.codegraph/config.json` exclude list. (3) Re-index via `scripts/codegraph_setup.sh` . (4) Validate via `scripts/codegraph_validate.sh` probe resolving a symbol that lives ONLY inside an own-org submodule. (5) Paired §1.1 mutation: temporarily add own-org path to exclude → validate FAILs → restore.

Composes with §11.4.74, §11.4.78, §11.4.10, §1.1, §107.

Canonical authority: `Constitution.md` §11.4.79.

Non-compliance is a process violation; severe cases (own-org submodules silently excluded without an audit trail in `.codegraph/config.json`) are release blockers.

§ 11.4.80 — CodeGraph regular-update + sync automation mandate

Direct user mandate (verbatim, 2026-05-21): “We MUST regularly check for the updates and execute codegraph npm updates so the latest version of it is always installed [...] Make sure we have proper full automation bash scripts which will run regularly and that these are part of the constitution Submodule [...] Make sure all updates, sync processes we do and important codegraph related events are all documented under docs/codegraph in Status and Status_Summary documents [...]”

Three deliverables in this constitution submodule: (1) `scripts/codegraph_update.sh` — npm-installs latest; §107 anti-bluff verifies `codegraph --version` reflects the new version. (2) `scripts/codegraph_sync.sh` — runs `status` → `sync` → `status` → project’s validate; appends to both project’s + constitution’s `docs/codegraph/Status.md` . (3) `docs/codegraph/Status.md` + `Status_Summary.md` ledgers, exported per §11.4.65 to `.html` + `.pdf` .

Cadence: weekly floor (§11.4.45). Scripts inherited by reference (§3) — Qwen Code consuming projects invoke them at `${CONST_DIR}/scripts/codegraph_*.sh` , never copy.

Composes with §11.4.78, §11.4.79, §11.4.10, §11.4.45, §11.4.65, §107, §1.1.

Canonical authority: `Constitution.md` §11.4.80.

Non-compliance is a process violation; severe cases (>2 weeks since last update + open AI-agent work) are release blockers.

§ . . – Cross-platform-parity mandate (User mandate, – –)

Direct user mandate (verbatim, 2026-05-21): “Any Linux-only blocker / issue we have MUST BE created macOS and other supported platforms equivalent! So, depending on platform proper implementation will be used for particular OS! EVERYTHING MUST BE PROPERLY EXTENDED AND UPDATED!”

Every consuming project whose supported-platforms manifest lists more than one OS MUST ship per-OS-equivalent implementations + tests for every feature/gate/challenge/mutation that depends on platform-specific primitives. Runtime dispatch via `uname -s` (or equivalent platform detection). Three sub-mandates: **(A)** Per-OS implementation REQUIRED — `cgroup/systemd/ /proc` primitives have documented equivalents (POSIX `setrlimit / ulimit`, `launchd`, `rctl`, Job Object). **(B)** Per-OS tests REQUIRED — every gate test has `case "$(uname -s)"` in branches with positive captured evidence per §11.4.2 + §11.4.5 in each branch; SKIP-with-reason only when the platform genuinely cannot enforce. **(C)** Honest kernel-gap citation + adjacent equivalent test REQUIRED where no equivalent exists (canonical: XNU does NOT enforce `RLIMIT_AS` for unprivileged processes → SKIP with exact reproducer + adjacent test of what IS enforced, e.g. `RLIMIT_CPU + SIGXCPU`). The adjacent test is itself anti-bluff per §11.4 with a paired §1.1 mutation.

Per-OS equivalence catalogue (canonical): `systemd-run --user --scope ↔ POSIX ulimit -t -u / launchd; cgroup MemoryMax ↔ XNU gap (use RLIMIT_CPU adjacent) / rctl / Job Object; cgroup TasksMax ↔ RLIMIT_NPROC ; /proc/<pid>/oom_score_adj ↔ no Darwin/BSD equivalent.`

Composes with §11.4.1 / §11.4.2 / §11.4.3 (strictened — SKIP only when kernel cannot) / §11.4.4 / §11.4.5 / §11.4.6 / §11.4.20 / §11.4.27 / §11.4.69 / §11.4.70 / §107. Pre-build gate `CM-CROSS-PLATFORM-PARITY` + paired §1.1 mutation. No escape hatch.

Canonical authority: `Constitution.md` §11.4.81.

Non-compliance is a release blocker on multi-platform projects.

§ . . – Iteration-speedup discipline mandate (User mandate, – –)

Direct user mandate (verbatim, 2026-05-22): “How can we speed-up this whole development and fixing process? ... all speed optimizations critical rules and mandatory constraints MUST BE all added into our root (constitution Submodule) Constitution.md, CLAUDE.md, AGENTS.md and QWEN.md!”

Iteration cycle time bounds the project’s defect-discovery rate. Slow cycles ship fewer validated features per unit of operator time. The 2026-05-22 forensic session witnessed ~3 hours of operator wait on a job whose intrinsic compute was ~90 min — every wasted minute came from a missing speedup discipline.

Every consuming project’s build / test / commit / debug pipeline MUST adopt the following 9 speedup disciplines AS MANDATORY:

- **(A) Phase 1 forensic before any speculative source patch.** Speculative patches without root-cause evidence are §11.4.6 + §11.4.82 violations.
- **(B) Live-ADB-First (or live-equivalent) before any rebuild.** §11.4.51 strengthened to release-blocker. Skipping live-probe for LIVE_ADB_TESTABLE changes = §11.4.82 violation.
- **(C) Pre-flight before rebuild orchestrators.** 30-sec readiness check verifies device + sink reachability, memory budget, lock files, orphan processes.
- **(D) Persistent build caches outside containers.** `ccache` + Soong / `sccache` / Gradle daemon state bind-mounted to host. Subsequent rebuilds drop from ~40 min to 5-15 min.
- **(E) Module-only rebuild for `CONFIG_*=m` driver patches.** `make modules` on single driver dir saves 15-20 min/cycle when applicable.
- **(F) Parallel multi-device testing.** Every owned validation device runs the autonomous cycle concurrently with separate output dirs. Catches per-device defects one cycle earlier.
- **(G) Subagent scope ≤30 min + worktree isolation + single-responsibility.** Empirical pattern: 8-task subagents stall, 2-3-task subagents complete.
- **(H) Lock-file + stale-process hygiene.** Clean `.git/index.lock` + orphan processes on session start. Disable auto git- `gc` in concurrent multi-agent repos.

- **(I) Cycle telemetry per §11.4.24.** Every cycle logs commit, per-phase wall-clock, speedup flag set, outcome. Weekly aggregation surfaces empirical ROI per discipline.

Composes with §11.4.4 / §11.4.6 / §11.4.9 / §11.4.20 / §11.4.24 / §11.4.42 / §11.4.43 / §11.4.50 / §11.4.51 / §11.4.52 / §11.4.58 / §11.4.70 / §12.7 / §107. Pre-build gate `CM-ITERATION-SPEEDUP-DISCIPLINE` (when implemented) + paired §1.1 mutation. No escape hatch — no `--skip-phase1` , `--no-pre-flight` , `--rebuild-everything` , `--unlimited-subagent-scope` , `--ignore-locks` , `--no-telemetry` flag exists.

Canonical authority: `Constitution.md` §11.4.82.

Non-compliance is a release blocker.

§ . . — docs/qa/ end-user evidence mandate (User mandate, — —)

Direct user mandate (verbatim, 2026-05-22): “every feature that ships MUST carry a recorded e2e communication transcript + any attached materials under `docs/qa/<run-id>/` (per-feature subdirectories). A feature with no QA transcript is itself a §107 PASS-bluff — it claims to work but has no auditable runtime evidence. Bot-driven automation MUST preserve full bidirectional communication threads as proof.”

Every feature that ships MUST carry a recorded end-to-end communication transcript plus attached materials committed under `docs/qa/<run-id>/` . Transcripts MUST be full bidirectional. Attached materials live in-repo (no external-only links — §11.4.13 sink-side violation). Bot-driven / agent-driven QA MUST preserve the full conversation thread as the proof artefact. CI release gates MUST refuse to tag a version whose feature-shipping commit lacks its `docs/qa/<run-id>/` .

Composes with §11.4.2, §11.4.5, §11.4.13, §11.4.65, §11.4.69, §107, §1.1.

Canonical authority: `Constitution.md` §11.4.83.

Non-compliance is a release blocker.

§ . . – Working-tree quiescence rule for subagent commits (User mandate, – –)

Short tag: `working-tree quiescence` .

Direct user mandate (verbatim, 2026-05-22): “no subagent commit may proceed while any concurrent mutation gate is in flight in the same checkout. Before `git add` , the committing agent MUST `grep` its own working tree for mutation markers (`MUTATED` for paired , `// always pass` , `return json.Marshal` shortcut paths, etc.). Any unexplained file in the staging area triggers ABORT.”

No subagent commit may proceed while any concurrent mutation gate is in flight in the same checkout. Pre- `git add` : `grep` for mutation markers (`MUTATED` for paired , `// always pass` , `return json.Marshal` shortcuts, `// MUTATION` annotations, `_mutated_*` filenames). Cross-check `git status --porcelain` against declared scope → unaccounted entries ABORT.

Lesson (forensic). A consuming project’s logo-fix subagent (Herald commit `72e81ab` , 2026-05-21) swept a `// always pass` JWT-bypass mutation residue into an unrelated commit, pushed to all four mirrors before being caught. Fix `d5bd360` landed within the hour, but the security-defect window is real and the lesson is constitutional.

Operative rule. (1) Pre- `git add` `grep` + scope-match → unaccounted ABORT. (2) Active mutation gates MUST be serialised before unrelated commits. (3) Concurrent subagents in same checkout MUST use lockfile (`.git/MUTATION_IN_PROGRESS`) or `git worktree add` . (4) Pre-push `mutation-residue-scanner` BLOCKS pushes containing mutation markers.

Composes with §1.1, §11.4.20, §11.4.70, §11.4.27, §11.4.10, §11.4.71, §107.

Canonical authority: `Constitution.md` §11.4.84.

Non-compliance is a release blocker.

§ . . – Stress + Chaos Test Mandate (User mandate, – –)

Short tag: `stress-chaos-mandate` .

Direct user mandate (verbatim, 2026-05-24): “Every fix or improvement you do MUST BE covered with full automation stress and chaos tests so we are sure nothing can break the functionality and all edge cases are monitored and polished and additionally fixed if that is needed! Everything must produce rock solid proofs and follow fully no-bluff policy!”

Every fix MUST ship with full-automation stress + chaos test suites. Happy-path-only coverage is a §11.4 / §107 PASS-bluff at the resilience layer.

Stress = sustained load ($N \geq 100$ iters OR ≥ 30 s) + concurrent contention ($N \geq 10$ parallel) + boundary conditions (empty/max/off-by-one). **Chaos** = failure-injection per §11.4.69 closed-set (process-kill, network-drop, input-corruption, OOM, disk-full, state-corruption) + verifiable recovery.

Anti-bluff: every stress + chaos PASS cites a captured-evidence artefact (latency.json / categorised_errors.txt / state_delta_snapshot.json) per §11.4.5 + §11.4.69. Helper library `stress_chaos.sh` provides `ab_stress_run` / `ab_stress_concurrent` / `ab_chaos_*` primitives. Chaos cleanup non-negotiable — corrupt-restore / disk-fill-cleanup / process-restart MUST run in trap EXIT.

4-layer per §11.4.4(b): pre-build gate (test files exist + parse) + paired meta-test mutation + on-device test (if LIVE_ADB_TESTABLE) + HelixQA Challenge entry. Composes with §11.4 / §11.4.1 / §11.4.5 / §11.4.6 / §11.4.43 / §11.4.50 / §11.4.52 / §11.4.69 / §11.4.83.

Canonical authority: `Constitution.md` §11.4.85.

11 4 86
Non-compliance is a release blocker. No `--skip-stress` , `--no-chaos` , `--happy-path-suffices` flag exists.
2026 05 25

§ . . — Roster/corpus-backed Status-doc auto-sync mandate (User mandate, — —)

Forensic anchor (verbatim): “Make sure that assets and players Status docs are ALWAYS regularly updated and in sync like all others Status docs — any time we add or modify the assets content(s) or we change or add new / remove existing pre-installed video and audio player apps! This MUST WORK OUT OF THE BOX!”

Status docs (§11.4.45) backed by a **tracked roster** (installed apps/components) or **asset corpus** (test/media directory) MUST resync out of the box the moment a member changes — Status doc + Status_Summary + HTML + PDF, mechanically. Mechanism (all must hold): (1) drift-proof **fingerprint** = sha256 of sorted member list (NOT mtime), persisted in a sidecar; (2) **sync helper** regenerates fingerprint + re-exports (via §11.4.65); (3) **pre-build gate** FAILs on fingerprint drift (mirrors §11.4.12 + §11.4.45); (4) paired §1.1 mutation. Composes with §11.4.12 / §11.4.45 / §11.4.53 / §11.4.56 / §11.4.57 / §11.4.59 / §11.4.60 / §11.4.65 / §11.4.6. Universal (§11.4.17) — consuming project supplies the specific docs/roster/helper/gate per §11.4.35.

Canonical authority: Constitution.md §11.4.86.

^{11 4 87} Non-compliance is a release blocker. No `--skip-roster-sync`,
`--allow-status-drift` flag exists.

^{2 0 2 6 0 5 2 6}

§ . . — Endless-loop autonomous work + zero-idle agent dispatch + anti-bluff testing mandate (User mandate, — —)

When operator instructs an AI agent to “continue in endless loop fully autonomously” (or semantically-equivalent), the agent MUST treat as HARD-CONTRACT covenant: (A) continue until docs/Issues.md non-terminal Status entries = 0 AND docs/CONTINUATION.md §3 Active work empty AND no subagent in-flight AND no external dep in-flight; (B) dispatch background subagents for parallelisable work — main + subagents concurrent, “waiting for results” is the ONLY idle reason; (C) every closure lands four-layer test coverage per §11.4.4(b) with captured-evidence “physical proofs” (audio/video/network/UI/sysfs) — metadata-only / config-only / absence-of-error / grep-without-runtime PASS are critical defects; (D) §11.4 anti-bluff covenant family operative end-to-end (tests AND HelixQA Challenges bound equally per forensic anchor “tests pass but features don’t work”); (E) loop terminates ONLY on all-conditions-met, explicit operator STOP, host-safety demand, or scheduled wake on known-future-actionable signal.

Composes with §11.4 / §11.4.1 / §11.4.2 / §11.4.4 / §11.4.5 / §11.4.6 / §11.4.7 / §11.4.20 / §11.4.27 / §11.4.42 / §11.4.43 / §11.4.50 / §11.4.52 / §11.4.58 / §11.4.68 / §11.4.69 / §11.4.70 / §11.4.83 / §11.4.85 / §11.4.86 / §12.10. Pre-build gate CM-COVENANT-114-87-PROPAGATION + paired §1.1 mutation.

Canonical authority: `Constitution.md` §11.4.87.

Non-compliance is a release blocker. No `--idle-OK` , `--skip-endless-loop` , `--bluff-permitted-for-this-task` , `--metadata-only-test-suffices` , `--no-physical-proof-required` flag exists.

Companion documents

- `Constitution.md` — authoritative universal Constitution. ALWAYS the tie-breaker.
- `CLAUDE.md` — Claude Code agent's view of the universal constitution.
- `AGENTS.md` — OpenCode / Cursor / generic-tooling view.
- `README.md` — high-level overview + multi-mirror contract.

When operating in a consuming project, ALSO read the project's local `QWEN.md` / `CLAUDE.md` / `AGENTS.md` for project-specific extensions; these MUST NOT weaken any universal rule but MAY add stricter project-specific constraints.

2026 05 26

§ . . — Background-push mandate:
commit-lock release immediately after commit,
push runs detached (User mandate,
- -)

Forensic anchor: 2026-05-26 `commit_all.sh` held flock ~5h on synchronous post-commit push. Mandate: (A) flock release IMMEDIATELY after commit; (B) push detached via `nohup` + `disown`; (C) `push_all.sh` per-remote flock — same-remote serializes, different-remote parallel; (D) failures → `qa-results/push_failures/` , autonomous loop checks per §11.4.87(A); (E) `--sync-push` escape for §11.4.41 force-push only. Composes §2.1 / §9.2 / §11.4.41 / §11.4.42 / §11.4.71 / §11.4.87. Gates `CM-COVENANT-114-88-PROPAGATION` + `CM-BACKGROUND-PUSH-WIRED` + paired §1.1 mutations.

Canonical authority: `Constitution.md` §11.4.88. Non-compliance is a release blocker.

§ . . – Background test execution mandate (User mandate, – –)

Forensic anchor 2026-05-27: conductor invoked long pre_build synchronously, blocking main stream 6-7 min on §JV/§JW/§JX/§JY scaffolding. Symmetric to §11.4.88 at test-execution layer. Mandate: (A) long tests (>30s) run via `nohup ... > <log> 2>&1 & + disown`; (B) main stream proceeds to §11.4.42 priority queue immediately; (C) hard-dependency gating via poll/pgrep — surface §11.4.66 options if still running; (D) failures land in log files; (E) foreground only <30s OR operator authorisation; (F) per-script flock. Composes §11.4.42 / §11.4.66 / §11.4.82 / §11.4.84 / §11.4.85 / §11.4.87 / §11.4.88. Gates `CM-COVENANT-114-89-PROPAGATION` + `CM-BACKGROUND-TEST-EXECUTION-WIRED` + paired §1.1 mutations.

^{11 4 90}
Canonical authority: `Constitution.md` §11.4.89. Non-compliance is a release blocker.
_{2 0 2 6 0 5 2 7}

§ . . – Obsolete status + obsolescence audit (User mandate, – –)

§11.4.15 Status closed-set + terminal `Obsolete` (→ `Fixed.md`) . Reasons: `superseded-by-design-change` | `superseded-by-later-mandate` | `feature-removed` | `duplicate-of` | `unsupported-topology` | `not-reproducible` (`not-reproducible` = a reported defect that does NOT reproduce on the canonical tree/baseline — an environment/isolated-worktree artifact, not a real product defect). `**Obsolete-Details:**` line mandatory (Since + Reason + Superseding-item + Triple-check evidence). Colorizer `cell-status-obsolete` = light-gray + strikethrough. Audit cadence per §11.4.40 + §11.4.42. Triple-check non-negotiable.

Canonical authority: `Constitution.md` §11.4.90. Release blocker.

§ . . – Summary-doc clarity (User mandate, – –)

Every summary entry's description: self-contained, meaningful, ≥ 6 words / ≥ 40 chars, SUBJECT + PROBLEM/GOAL. Anti-patterns forbidden: section labels, bare metadata, §-letter alone. Generators extract from H1/H2 heading. Pre-build gate

CM-SUMMARY-CLARITY-DESCRIPTIONS .

Canonical authority: 2026 05 27 Constitution.md §11.4.91. Release blocker.

§ . . – Multi-pass change-evaluation (User mandate, – –)

5-pass evaluation BEFORE commit-ready: (1) Main-task verification; (2) Regression-blast-radius; (3) Cross-feature interaction; (4) Deep-research per §11.4.8; (5) Anti-bluff confirmation. Doc per pass. Trivial exemption: typo/revision/MD-export ONLY if zero source touched.

Canonical authority: 2026 05 27 Constitution.md §11.4.92. Release blocker.

§ . . – SQLite-SSoT for workable items (User mandate, – –)

docs/.workable_items.db SQLite — DB authoritative, MD derived. Schema: items + item_history + obsolete_details + operator_block_details + firebase_metadata + meta. Go binary cmd/workable-items/ sync md ↔ db + diff + validate + add + close. Bidirectional regen byte-identical. Anti-bluff: unit+integration+stress+chaos per §11.4.85. Go binary in constitution submodule per §11.4.74. 6-phase migration §LA.

Canonical authority: Constitution.md §11.4.93. Release blocker.

§ . . – Zero-idle priority-first parallel-by-default (User mandate, – –)

Binds §11.4.20/42/58/70/72/82/87/88/89 — always-on. Idle ONLY: externally blocked, operator STOP, §12 host-safety. Before any wake/sleep survey parallel-work queue. Priority MANDATORY (§11.4.72 audio first). Subagent-driven default non-trivial. Background default >30s. Stability per §11.4.92 + §11.4.84 + §12.6-9. Updates at milestones.

Canonical authority: ^{2026 05 27} Constitution.md §11.4.94. Release blocker.

§ . . – Workable-items DB TRACKED in git (User mandate, – –)

§11.4.93 amendment: docs/workable_items.db TRACKED, NEVER gitignored. Authoritative source. Every mutation → stage+commit+push DB + MD. WAL-checkpoint before commit. §11.4.77 does NOT apply. Pre-build gates CM-COVENANT-114-95-PROPAGATION + CM-WORKABLE-ITEMS-DB-TRACKED .

Canonical authority: ^{2026 05 27} Constitution.md §11.4.95. Release blocker.

§ . . – Safe-parallel-work-with-long-build (User mandate, – –)

SAFE during AOSP build: (A) docs/MD; (B) scripts/; (C) pre-build gates; (D) on-device tests; (E) constitution edits+push; (F) submodule push per §11.4.88; (G) read-only ADB probes; (H) subagent dispatch; (I) web research; (J) workable-items DB ops; (K) pre-build + meta-test execution backgrounded.

UNSAFE: (α) git checkout/reset on source; (β) mass deletes/renames under built source; (γ) APK-built submodule pointer; (δ) out/; (ε) make clean; (ζ) container destruction; (η) disk-fill; (θ) §12 host-safety.

Conductor dispatches ≥1 (A)-(K) per pause. “Build running, nothing else to do” NEVER true.

Canonical authority: ^{2026 05 27} Constitution.md §11.4.96. Release blocker.

§ . . – Max-use-of-idle + progress-update cadence (User mandate, – –)

1. Every minute progressable idle = §11.4.97 violation, dispatch CONTINUOUSLY; (B) 1-line operator updates at commit/subagent/anchor/evidence/migration — no prompt; (C) continuous physical-proof gathering per §11.4.5/6/69; (D) composes with operating-mode anchor family; (E) idle-only-when-blocked closed-set unchanged from §11.4.94(A). Pre-build gates

1.1 CM-COVENANT-114-97-PROPAGATION + CM-IDLE-TIME-AUDIT .

Canonical authority: ^{2026 05 28} Constitution.md §11.4.97. Release blocker.

§ . . – Full-Automation Anti-Bluff (User mandate, – –)

Forensic anchor: “Make sure we have full automation testing of all scenarios with real bot, main group and users without any manual intervention or contribution of real user! ... No bluff is allowed!”

Closes the manual-intervention gap §11.4+85+87+89+94 didn’t forbid. (A) every test MUST self-drive end-to-end — PASS/FAIL/SKIP-with-reason without further human action; (B) single exception = one-time credential bootstrap OUTSIDE test execution (.env, ~/.bashrc, OAuth, MTPROTO-session-activation); (C) concrete live requirements: (1) no “operator MUST type” prompts — drive via MTPROTO/real-user-API/webhook-fixture/in-process loopback; (2) no UUIDs colliding with active dev session (Herald 2026-05-28: silent exit -1 lesson); (3) no 60s human-response windows (§11.4.50 violation); (4) re-runnability -count=3 with self-cleaning state; (5) audit COMPLIANT vs NON-COMPLIANT; (6) no false-positive PASS — silent-skip-as-PASS + stale-evidence forbidden, §11.4.3 SKIP-with-reason correct; (D) composes §11.4.85+89+87+94 = continuously-validated fully-automated non-flake anti-bluff regime; (E) inheritance per §11.4.35 — restate literal 11.4.98 in every consumer; pre-build gate CM-COVENANT-114-98-PROPAGATION ; paired §1.1 mutations strip → FAIL; (F) enforcement: manual-action commit BLOCKED; NON-COMPLIANT after 30 days → §11.4.90 Obsolete citing §11.4.98.

Canonical authority: Constitution.md §11.4.98. Release blocker.

§ . . — Latest-Source Documentation Cross-Reference (User mandate, – –)

Forensic: “ALWAYS check against latest versions of services we use web/online docs before creating instructions! ... mandatory rules / constraints, result is consistency and safety of created instructions, guides and manuals!”

Case study (Herald 2026-05-28): MTPProto guide recommended VoIP + omitted `recover@telegram.org` pre-login email; contradicted official Telegram + gotd/td docs; could have caused permanent ban. Misguidance-by-stale-docs = §11.4 PASS-bluff at documentation layer.

1. Pre-commit cross-reference: fetch LATEST official online docs (WebFetch/MCP/ direct browsing — NEVER training data) for every operator-facing instruction doc; cite source URL+date in “## Sources verified” footer AND commit-message footer; seek secondary authoritative sources for sparse official docs. (B) Negative findings documented explicitly. (C) STALE after 6 months (90 days for risk-classified services); re-verify at vN.0.0, on breaking-change, on operator-error. (D) Risk-classified: messengers/cloud/payment/AI/code-hosting/package-managers. (E) Composes with §11.4.92 Pass 4 INDEPENDENT — cannot substitute. (F) Inheritance per §11.4.35 — restate literal `11.4.99` in every consumer’s CLAUDE/AGENTS/QWEN. (G) Enforcement: missing-footer BLOCKED; stale-beyond-grace → §11.4.90 Obsolete.

Canonical authority: `Constitution.md` §11.4.99. Release blocker.

§11.4.100 — RETIRED. Demoted to consumer project (a consuming project’s video-color/visual-quality fidelity) per §11.4.17/§11.4.35 — project-specific (RK3588/MPV/Arvus), not universal. See the consuming project’s Constitution/CLAUDE/AGENTS/QWEN.

§ . . – Autonomous-decision-over-blocking (User mandate, – –)

Forensic: “when working in endless working loop fully autonomously try to decide most properly about points which would block execution and wait for us. If we haven’t answered now work would be blocked whole night! If possible and if that will not cause any issues make proper and most reliable and safe decision so we achieve maximal efficiency and work gets fully done!”

In autonomous / endless-loop mode (per §11.4.87) the agent MUST minimize operator-blocking and make the safe, reliable, reversible decision itself — work NOT stalled overnight waiting for input. §11.4.87 = keep working; §11.4.101 = HOW to clear the decision points.

Decision rule (proceed autonomously when ALL): (a) reversible OR pre-op backup per §9.2; (b) safe choice determinable from captured evidence per §11.4.6 (LIKELY ≠ determination); (c) wrong-choice blast radius bounded AND recoverable; (d) composes with §11.4 + §12 + §9. Block-only-when (BLOCK via §11.4.66 ONLY when ALL): irreversible AND high-blast-radius AND choice undeterminable from evidence — external-account state, inaccessible hardware, destructive-op-without-backup, force-push (§9.2 + §11.4.41), spending / third-party send. Operator-blocked per §11.4.21 only after this fires + self-resolution-exhaustion audit. Maximize-progress-while-blocked: a block parks one work unit, not the loop — keep progressing NON-blocked items per §11.4.87 + §11.4.94; pose-and-idle = §11.4.94 + §11.4.97 violation. Composes §11.4.6/21/40/41/66/87/94/§9.2/§12. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-101-PROPAGATION (literal 11.4.101) + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: Constitution.md §11.4.101. Release blocker. No -- always-block-on-decision / --never-decide-autonomously / --skip-decision-rule / --block-without-self-resolution escape.

§ . . – Systematic-debugging always-on + always-loaded skill-discovery + plugin availability (User mandate, – –)

Forensic: “ALWAYS trigger / start the”/superpowers:systematic-debugging” skills when any issues happen! ... we MUST activate the skill(s) and make strongest efforts in full in depth analysis / debugging and determine root causes ... “/using-superpowers” skill is ALWAYS loaded, applied and used! All dependencies (plugins) that Claude Code or other market places are offering MUST BE installed if these are not already available for loading and use!”

1. On ANY spotted issue/bug/test-failure/gate-failure/regression/misalignment/inconsistency/unexpected-behaviour the agent MUST activate

`superpowers:systematic-debugging` (or platform-equivalent) BEFORE proposing/writing/applying any fix — Iron Law: NO FIXES WITHOUT ROOT CAUSE INVESTIGATION FIRST. Four-phase arc: root-cause → pattern → hypothesis (prove/disprove with captured evidence) → implementation (against proven cause only). Guess-and-retry / symptom-patch / re-run-hoping-it-passes (“probably transient/flaky”) without completed investigation = §11.4.102 violation; calling a failure `transient / flaky / intermittent` without captured forensic evidence = simultaneous §11.4.6 + §11.4.7 violation. (B) `superpowers:using-superpowers` (or platform-equivalent skill-discovery) ALWAYS loaded + applied at session start, consulted before any task — if ANY skill could apply (even 1%) it MUST be invoked, never improvised. (C) Every mandated skill plugin/dependency (Claude Code marketplaces or other runtime ecosystems) MUST be installed + loadable BEFORE dependent work; missing plugin blocking a mandated skill = release-blocker. Install: runtime’s own path — Claude Code `/plugin` flow (add marketplace → install → confirm skill in available-skills list); other runtimes’ documented installer. Anti-bluff: confirm via live capability list (install exit 0 ≠ skill loadable, §11.4.80 lesson). (D) AUTOMATIC + STANDING DEFAULT (User mandate 2026-07-13, verbatim “Make sure we automatically turn on the systematic-debugging plugin and skills every time we work on any issue! Add this as MANDATORY RULE to our root constitution Submodule and all of its relevant files!”): clause-(A) activation is AUTOMATIC — auto-activate `superpowers:systematic-debugging` on ANY spotted issue/bug/test-failure/gate-failure/regression/inconsistency/unexpected-behaviour WITHOUT an operator prompt/request/per-issue confirmation, the

MOMENT an issue is detected (composes §11.4.126 from-first-prompt default); plugin + skills confirmed installed+loadable per (C) before dependent work (§11.4.6, never assumed); “did not activate because not asked” = §11.4.102(D) violation; STRENGTHENS (A) never weakens it — (A)/(B)/(C) + four-phase-arc + Iron-Law intact. Composes §11.4.4/6/7/8/43/70/82(A)/92/126. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-102-PROPAGATION (literal 11.4.102) + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: Constitution.md §11.4.102. Release blocker. No

```
--skip-systematic-debugging / --guess-and-retry-OK /  
1 1 4 1 0 3  
--symptom-patch-permitted / --skip-skill-discovery / --plugin-  
optional / --missing-plugin-is-warning escape.  
2 0 2 6 0 5 2 9
```

§ . . – Continuous parallel-stream working routine (User mandate, – –)

Forensic: “Do this working approach continuously and make it part of regular working routine and add it to the Constitution Submodule documented fully.”

Promotes the proven multi-stream pattern into the project’s standing default working routine (not a per-request opt-in). Load-bearing invariant: main work stream MUST always stay FREE. (A) Main stream FREE — ALL commit AND push run detached (nohup ... & + disown per §11.4.88), never blocking on push or slow mirror. (B) ≥3 parallel background streams at all times + auto-backfill (User mandate 2026-05-29; raised from ≥2) — at least three subagent-driven streams (§11.4.70/§11.4.20, isolated §11.4.58 + §11.4.84) alongside main whenever three-plus non-contending actionable items exist; the moment any one stream is FULLY done a new stream MUST immediately start + take its place (next-highest-priority non-contending item), count NEVER drops below 3 while actionable items remain. Idle below 3 only when no remaining items OR all externally blocked (§11.4.94/97/101). Band 3–6, bounded §12.6 60% mem + §11.4.58 6-agent cap. (C) Most-critical + most-visible first; audio always top §11.4.72 + §11.4.42. (D) Safe-during-build scope only — while heavy gradle / m -j /42 GB containerised AOSP build (§12.9) runs, streams restrict to §11.4.96 SAFE catalogue; NEVER a second concurrent heavy build (§12.8). (E) Heavy anti-bluff every closure (§11.4.5/6/50/69/102; root cause proven or UNCONFIRMED: / UNKNOWN: / PENDING_FORENSICS:). (F) Idle ONLY when genuinely externally blocked

(hardware/network/in-flight build-test-push) OR operator STOP OR §12 host-safety per §11.4.94(A)+§11.4.97; block parks one work unit not the loop (§11.4.101). Composes §11.4.58/70/72/87/88/89/94/96/97/101/102/42/84/§12.6-§12.9/§9.2. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-103-PROPAGATION (literal 11.4.103) + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: Constitution.md §11.4.103. Release blocker. No --block-main-stream / --synchronous-commit / --synchronous-push / --single-stream-only / --skip-parallel-streams / --serialise-actionable-work / --idle-without-queue-survey escape.

§ . . – Participant identity, attribution & notification-tagging (User mandate, – –)

Forensic: “Every supported messenger must relate messages to participants (Subscribers/Users); the same logical person may have a different username on every messenger. Workable items must carry who created them and who they are assigned to. Notifications must @-tag the right participant — but never the operator (who drives the system) and never the system agent.”

MANDATORY for every consumer shipping a messenger/notification surface (Herald + flavor binaries = reference impl; others inherit §11.4.35). Detailed spec: Herald docs/design/PARTICIPANT_ATTRIBUTION.md — restated, not redefined. (A) Every messenger MUST relate messages to a **Participant** (logical Subscriber/User); SAME person MAY have a DIFFERENT username per messenger (logical subscriber: canonical messenger-neutral handle , kind ∈ {human, agent, service} ; + per-channel aliases channel / channel_user_id / @username for tagging). Canonical handle closed set: Claude (reserved system-agent sentinel; never tagged) OR a subscriber @username . (B) Operator = env var HERALD_<CHANNEL>_OPERATOR_USERNAME , not a DB flag — the one human who drives via the agent CLI; a normal Participant whose handle equals that value. (C) Workable items MUST carry created_by + assigned_to (canonical handles): CLI-prompt → Operator; system/agent-detected → Claude ; received-message → sender’s resolved @username ; assigned_to defaults to Operator, overridable; legacy items carry "" and MUST still parse + validate. (D) Tagging matrix: tag assigned_to if human ≠ Operator; tag created_by if human ≠

Operator ≠ `claude` ; NEVER tag `claude` or the Operator; de-dup; resolve to channel `@username` , skip if absent there. (E) Anti-bluff (composes §11.4): real SQLite round-trip with the new columns byte-identical (incl. legacy fixtures WITHOUT the fields); tagging matrix proven by a truth-table + cell-flip mutation forcing FAIL; E2E real-event → real-message asserting the exact `@username` s + a NEGATIVE case proving the Operator is NOT tagged; evidence under `docs/qa/<run-id>/` . Composes §11.4 + §11.4.1..16 / §11.4.5 / §11.4.69 / §11.4.50 / §11.4.91 / §11.4.93 / §11.4.95 / §1.1. Classification: universal (§11.4.17); no-messenger projects inherit latently (§11.4.96 pattern). Propagation gate `CM-COVENANT-114-104-PROPAGATION` (literal `11.4.104`) + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: `Constitution.md` §11.4.104; detailed spec Herald `docs/design/PARTICIPANT_ATTRIBUTION.md` . Release blocker. No `--skip-attribution` / `--no-participant-tagging` / `--tag-operator-anyway` / `--attribution-later` escape.

§ . . – Natural-language intent recognition & clarification (User mandate, – –)

Forensic: “Users must NOT need to know command syntax (no `COMMAND: ...` prefix). They send a clear natural-language message; the System determines the intent. The System recognizes the commands it has; if none matches it infers the exact intent; if it is totally unable it replies, tags the user (`@user ...`), and asks to clarify precisely. We MUST always do our best to determine exact intent so we never annoy end users. This is a CORE part of the System.”

MANDATORY for every consumer shipping a messenger/command surface (Herald + flavor binaries = reference impl; others inherit §11.4.35). Detailed spec: Herald `docs/design/INTENT_RECOGNITION.md` — restated, not redefined. (A) No required command syntax — users MUST NOT need to know any syntax (no `COMMAND: prefix`); they send plain natural language, the System determines the intent. (B) Three-tier resolution, first that succeeds wins: TIER 1 recognize the System’s existing command set from natural language → action (confident deterministic match); TIER 2 when no command matches, infer the exact intent (LLM maps language → action), NEVER guessing; TIER 3 when neither a command nor a confident intent is determinable, REPLY to the message, TAG the

sender (@username , via the §11.4.104 IdentityResolver) and ask a PRECISE clarifying question NAMING the candidate intents — no guessing, no silent drop. (C) Never guess, never drop — a wrong action is worse than a clarifying question (composes §11.4.6); a message is NEVER silently dropped; only genuine ambiguity reaches Tier 3, which always replies-tags-and-asks. (D) Anti-bluff (composes §11.4): every tier ships unit + integration + E2E + full-automation tests with real captured evidence — Tier 1 truth-table (natural-language → action+fields + conservative negatives that MUST fall through to “no match”); Tier 3 E2E whose recorded reply body is EXACTLY @<sender> <specific question> + a NEGATIVE proving a clear command does NOT trigger clarify; a paired §1.1 mutation breaking the confidence guard (false-match) OR dropping the clarify tag MUST FAIL a test; evidence under docs/qa/<run-id>/ . Composes §11.4 + §11.4.1..16 / §11.4.6 / §11.4.104 / §11.4.5 / §11.4.69 / §11.4.98 / §1.1. Classification: universal (§11.4.17); no-messenger projects inherit latently (§11.4.96 pattern). Propagation gate CM-COVENANT-114-105-PROPAGATION (literal 11.4.105) + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: Constitution.md §11.4.105; detailed spec Herald docs/design/INTENT_RECOGNITION.md . Release blocker. No `--require-command-syntax` / `--guess-intent-ok` / `--skip-clarify` / `--drop-on-ambiguous` escape.
11 4 105
2026 05 31

§ . . — Docs Chain — mechanical documentation/DB sync engine (Operator mandate, — —)

Forensic: Docs Chain is the canonical mechanical enforcer of the documentation-sync mandates; consumers MUST use the engine instead of ad-hoc per-project doc-sync scripts, register chains via per-context YAML, and never accept a faked transform.

MANDATORY for every consumer. Docs Chain (the vasic-digital/docs_chain engine — a universal Go bidirectional document-and-database dependency-propagation engine) is the canonical mechanical enforcer of the documentation-sync mandates. Detailed spec: engine docs ~/Projects/docs_chain → docs/CONSTITUTION_INTEGRATION.md (distribution + inheritance + anchor mapping table) + docs/USE_CASE_CATALOGUE.md (chain recipes) — restated, not redefined. (A) Use the engine, never ad-hoc scripts — consumed by reference (the flat-layout

sibling `~/Projects/docs_chain /` the constitution-exposed path), inherited like §11.4.80's `codegraph_*` scripts: referenced, NEVER copied; ad-hoc `sync_*` / `generate_*` / `summary` / `update_readme_doc_links` scripts superseded + retired per registered context. (B) Consumer-owned contexts — the engine is project-agnostic; the consumer registers chains as data via `.docs_chain/contexts/*.yaml` (§11.4.28 decoupling); `state.json` + `*.docs_chain.tmp` gitignored. (C) Anchors mechanized (per CONSTITUTION_INTEGRATION mapping table): §11.4.12/.53/.45/.56/.57/.59/.60/.65/.86/.93/.95/.44 + §12.10 — content-hash change detection (NOT mtime, §11.4.86), atomic-rename + SQLite-txn commit + rollback (§9.2), both-dirty `sync` → conflict-not-silent-merge (§11.4.6), `verify` as the deterministic CI/pre-build gate (§11.4.50), per-run captured evidence to `qa-results/docs_chain/<run-id>/` (§11.4.69). (D) NOT a replacement for authoring discipline — the source author still writes the §11.4.44 revision header; the engine only keeps exports in sync. (E) Anti-bluff (composes §11.4): NEVER fakes a transform — a missing pandoc/weasyprint surfaces a typed `ToolAbsentError` + honest §11.4.3 SKIP-with-reason, never a fake PASS / partial write; every `sync` / `verify` carries real captured evidence. Composes §11.4 + §11.4.1..16 / §11.4.6 / §11.4.28 / §11.4.80 / §9.2 / §11.4.50 / §11.4.69 / §11.4.5 / §1.1 + the mechanized sync anchors. Classification: universal (§11.4.17); no derived-export/DB-sync projects inherit latently (§11.4.96 pattern). Status (§11.4.6): engine Phases 1–3 AND the CLI/YAML loader (Phase 4) IMPLEMENTED+tested — `cmd/docs_chain/main.go` registers `sync` / `verify` / `diff` (alias) / `doctor` , `go test -count=1 ./...` GREEN 7/7 packages; this CLI/loader is in the working tree but UNTRACKED and the engine is NOT yet a registered git submodule (in-tree at `./docs_chain` , HEAD = parent HEAD); only submodule distribution (Phase 6) remains PLANNED + OPERATOR-GATED. Propagation gate `CM-COVENANT-114-106-PROPAGATION` (literal `11.4.106`) + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: `Constitution.md` §11.4.106; detailed spec Docs Chain engine docs `docs/CONSTITUTION_INTEGRATION.md` + `docs/USE_CASE_CATALOGUE.md` (`vasic-digital/docs_chain`). Release blocker. No `--ad-hoc-sync-ok` / `--skip-docs-chain` / `--fake-transform` / `--sync-evidence-optional` escape.

§ . . – Anti-bluff AV/test-validation techniques mandate (User-driven research, – –)

Forensic (genericised, 2026-06-02): a test PASSEd on a SINGLE captured frame showing “a picture” — but it was a FROZEN/STALE frame from previously-played content (stale-producer/stuck-decoder), feature broken for the user while green; a sibling FLASHED media on the WRONG output for ~1 s before routing with no test sampling that window; a third shipped a comparator that PASSEd its own deliberately-degraded fixture (the analyzer was the bluff). §11.4.5 mandates captured evidence + a presence pass; §11.4.107 raises it to liveness + correct-routing + self-validated-analyzer.

Every test asserting audio/video output is genuinely playing/advancing MUST satisfy ALL: (1) a single captured frame is NOT proof — prove LIVE, ADVANCING frames over a window via a freeze-detection oracle (near-duplicate/
`freezedetect` -class OR perceptual-hash adjacent-frame distance, NOT byte-identical — byte-identity only a pre-filter); (2) an independent frame-advance counter from the platform’s compositor/decoder telemetry must increase (different-domain oracle — flat $\Rightarrow$ stuck decoder $\Rightarrow$ FAIL even if pixels appear to move); (3) loading/buffering is a distinct state — wait for genuine playback-start, never false-FAIL a still-loading stream nor false-PASS a spinner (timeout+unreachable $\Rightarrow$ SKIP §11.4.3, timeout+reachable $\Rightarrow$ FAIL); (4) not-stale-from-previous cross-check (new first frame $\neq$ previous last frame); (5) measured FPS / no-lost-frames within tolerance; (6) no-flash-on-the-wrong-output — high-frequency sampling of the non-target output during routing transitions, any content frame there $\Rightarrow$ FAIL (content-protection regime classified, never guessed); (7) drive through the realistic feed/UI path, not deep-link shortcuts (UI-not-introspectable $\Rightarrow$ operator-attended §11.4.52, never fake-PASS); (8) metamorphic relations for the oracle problem on un-owned streams (same content output-A vs output-B must match; paused $\Rightarrow$ counter stops; 2 $\times$ speed $\Rightarrow$ ~2 $\times$ advance); (9) full-reference quality metrics (SSIM/VMAF/ ΔE_{2000}) vs golden source for owned content; (10) mutation-test every analyzer with a golden-good + golden-bad fixture pair so the analyzer cannot bluff (PASS on golden-bad = bluff gate — §1.1 applied to analyzers); (11) per-channel audio RMS/loudness (EBU R128) + XRUN census (aggregate RMS misses a dead channel); (12) OCR overlay/subtitle needs a per-word confidence floor + ROI (avoids both

false-FAIL and false-PASS); (13) thresholds calibrated on the project's own fixtures, not hardcoded from literature (§11.4.6). Honest gap: true photon FPS at the sink has no clean software oracle — flag it, never claim it.

4-layer §11.4.4(b): pre-build gate (CM-AV-LIVENESS-NO-FROZEN-FRAME -class + analyzer-self-validation gate wiring golden-good/golden-bad fixtures into meta-test) + runtime/on-device test + paired §1.1 mutation (single-frame-only → FAILs; analyzer PASSing its golden-bad fixture → self-validation FAILs) + HelixQA Challenge. Every PASS via `ab_pass_with_evidence` citing motion / not-stale / frame-advance / fps / loudness / metamorphic / self-validation artefacts (`video_display` / `audio_output` / `subtitle_render` §11.4.69) — never a single screenshot. Classification: universal (§11.4.17) — platform-neutral, reusable by ANY project validating media playback or any pixel/audio output; project supplies its capture mechanism + calibrated thresholds per §11.4.35. Composes §11.4.5/.6/.50/.68/.69/.85 + §11.4.2/.3/.13/.48/.52/§1.1. Propagation gate `CM-COVENANT-114-107-PROPAGATION` (literal `11.4.107`) + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: `Constitution.md` §11.4.107. Release blocker. No `--single-frame-proves-playback` / `--skip-liveness` / `--byte-identical-freeze-OK` / `--no-frame-advance-counter` / `--allow-wrong-output-flash` / `--deep-link-shortcut-OK` / `--unvalidated-analyzer-OK` / `--aggregate-rms-suffices` / `--hardcoded-thresholds-OK` escape.

§ . . — Four-layer fix-verification + runtime-signature-as-definition-of-done mandate (systematic-debugging Phase . , - -)

Forensic anchor (genericised, 2026-06-03): across one batch multiple fixes were “green” at every gate yet NOT working for the end user — a source-committed change passed by both the source pre-build gate AND the post-build gate never reached the boot-time command-line embedded in the deployed boot artifact (feature dead despite all-green); sibling fixes correctly built into the system image were masked by stale per-user overlay copies from a prior deployment (running code = stale shadow, not the fresh build); deployment did not wipe the mutable overlay so the shadow survived; validation ran against the stale shadow and reported green on code never exercised. Per `superpowers:systematic-`

debugging Phase 4.5: each fix revealing a fresh “fixed-but-not-working” in a DIFFERENT place is ONE architectural VERIFICATION flaw, not independent bugs — patching symptoms is thrashing. A fix crosses FOUR distinct layers, and “fixed” at one does NOT imply fixed at the next: (1) SOURCE (committed — grep-the-source gate), (2) ARTIFACT (the change's BYTES landed in the produced artifact — image / bundle / installer / embedded command-line), (3) RUNTIME-ON-CLEAN-TARGET (active on a CLEAN/fresh deployment, the layer the end user experiences, no stale overlay shadowing it), (4) USER-VISIBLE (works for the end user — §11.4.5/§11.4.69 captured-evidence). The mandate (ALL hold): (1) **runtime-signature-as-definition-of-done** — a fix is DONE only when its declared runtime signature verifies on a CLEAN deployment; source-committed ≠ artifact-contains-it ≠ active-on-clean-target ≠ user-visible-working; (2) every fix declares ONE machine-checkable runtime signature (running-system property / sink-side report per §11.4.13 / §11.4.5/§11.4.69 captured-evidence / counter delta — NEVER a re-grep of the source); this registry of per-fix runtime signatures is the SINGLE SOURCE OF TRUTH for “fixed”, REPLACING “a gate greps the source”; (3) gates span all four layers — source (pre-build), artifact (post-build — assert the BYTES landed, not merely that the source still contains the change), runtime-on-clean-target (post-deploy — assert the runtime signature), user-visible (§11.4.5/§11.4.69); (4) eliminate the stale-deployment / shadow layer by construction — deployment MUST yield a CLEAN state (wipe the mutable overlay) OR a pre-validation assertion MUST prove `running-artifact == built-artifact` before any validation; validation against possibly-stale state is INVALID and its PASS is a §11.4 PASS-bluff; (5) meta-rule — ≥3 “fixed-but-not-working” discoveries in one cycle signal an architectural VERIFICATION flaw, not three independent bugs; on the 3rd STOP patching symptoms (§11.4.4), fix the VERIFICATION pipeline, re-certify EVERY item through it on a clean target; (6) a batch is “validated” only after COMPREHENSIVE per-item runtime-signature verification on a clean baseline, NOT after the touched items' own tests pass. Classification: universal (§11.4.17) — platform-neutral verification-pipeline disciplines reusable by ANY project that builds → deploys → validates an artifact; the project supplies its artifact format, clean-deployment / equal-artifact mechanism, and per-fix runtime-signature observables per §11.4.35. Composes §11.4.1/.2/.4/.5/.6/.27/.40/.46/.50/.52/.69/.102. Propagation gate `CM-COVENANT-114-108-PROPAGATION` (literal `11.4.108`) + recommended per-fix gate `CM-RUNTIME-SIGNATURE-REGISTRY` + paired §1.1 mutations (gate-code = separate work item).

Canonical authority: Constitution.md §11.4.108. Release blocker. No `--source-green-is-done / --skip-artifact-byte-check / --validate-against-running-state / --no-clean-deployment / --skip-runtime-signature / --spot-validate-touched-only` escape.

§ . . — Mandatory Anti-Forgetting Enforcement: PreToolUse Guard Hook + Subagent Constitutional Preamble + Orchestrator Pre-Action Checklist (Operator mandate)

UNCONFIRMED forensic anchor (pending operator's verbatim mandate quote).

Background: emulator subagents ran raw host-direct `emulator / adb` because the orchestrator forgot to inject the Containers-submodule rule at dispatch. A forgotten rule is not enforcement. Two-layer fix: (A) portable bash/awk

`PreToolUse hook ( constitution/scripts/hooks/guard-forbidden-commands.sh , no jq , no project-specific paths)` blocks host-direct emulator / force-push bypass / sudo / host-power at the tool-call boundary regardless of agent memory (the floor); (B) `constitution/docs/AGENT_GUARDRAILS.md` — canonical preamble the orchestrator pastes verbatim into every subagent dispatch + Orchestrator Pre-Action Checklist (the ceiling). Consuming projects wire the hook in `.claude/settings.json` or equivalent; maintain `docs/AGENT_GUARDRAILS.md` with `SUBAGENT CONSTITUTIONAL PREAMBLE` , `ORCHESTRATOR PRE-ACTION CHECKLIST` , and the anchor literal `11.4.109` ; provide a ≥ 20 -case hermetic hook test. The hook is inherited by reference — NEVER copied locally. Gates: `CM-ANTI-FORGETTING-ENFORCEMENT` + `CM-COVENANT-114-109-PROPAGATION` + paired §1.1 mutations (gate-code = separate work item). Classification: universal (§11.4.17). Composes §11.4.6/.10/.75/.76/.78/.79/.80/.81/.84/.98/.102/§12.

Canonical authority: Constitution.md §11.4.109; ref impl `constitution/scripts/hooks/guard-forbidden-commands.sh` + `constitution/docs/AGENT_GUARDRAILS.md` . Release blocker. No `--skip-pretooluse-hook / --no-guardrails-doc / --anti-forgetting-optional / --single-layer-sufficient` escape.

§ . . – Pre-build build-readiness verdict + change-impact clash detection mandate (operator mandate, – –)

Forensic anchor (genericised, 2026-06-03): a fix shipped a new system-property *read* but introduced no matching security-policy grant + no property-context type entry — the read was silently denied at runtime + the feature was dead, while every pre-build gate stayed green because the gate only grepped the source file.

Generalises: a change introduces a new dependency on a *second* artifact (security policy, context file, service registry, interface freeze-snapshot, symbol table, build-graph node) that the pre-build never cross-checks. This is §11.4.108's

SOURCE → ARTIFACT gap shifted LEFT to pre-build — most such clashes are statically catchable from the change diff itself. The mandate (ALL hold): (1) a single deterministic **READY-FOR-BUILD verdict** gates the rebuild (orchestrator refuses to start unless READY); (2) a **diff-driven change-impact + clash detector** cross-checks every newly-introduced second-artifact dependency (new property read ⇔ property-context type + read-grant; new service ⇔ service-context; new init service ⇔ security label; new/changed stable interface ⇔ freeze-snapshot updated; new native-lib dep ⇔ module/prebuilt resolves; new policy rule ⇔ every type/attribute defined; two batch changes on the same seam ⇔ collision acknowledged); (3) **coverage-completeness is a gate** — every changed file maps to ≥1 gate + ≥1 deployed-target test + ≥1 paired §1.1 mutation, baseline ratchets per §11.4.50; (4) **two-speed honesty** — grep-speed always-on gates vs REQUIRES_BUILD heavy gates (build-graph parse-only dry-run, full neverallow compilation, ABI diff) as diff-gated opt-in stages bounded per §12.6/§12.7; (5) every gate + wired analyzer is anti-bluff by paired §1.1 mutation; (6) **honest boundary** — a READY verdict proves static internal-consistency + ready-to-build, NOT that the feature works on the deployed target; the regime empties the *preventable* class, not the *all-defects* class (runtime/USER-VISIBLE remains §11.4.108's job). Classification: universal (§11.4.17) — platform-neutral pre-build-rigor disciplines reusable by ANY artifact-building project; the project supplies its property/service/policy registries, build-graph parse-only command, interface-freeze mechanism + changed-file → gate mapping per §11.4.35. Composes §11.4.1/.4/.6/.9/.27/.50/.67/.75/.92/.108.

Propagation gate CM-COVENANT-114-110-PROPAGATION (literal 11.4.110) + recommended CM-READY-FOR-BUILD-VERDICT / CM-CHANGE-IMPACT-CLASH-

DETECTOR / CM-COVERAGE-COMPLETENESS-GATE /

CM-BUILDGRAPH-DRYRUN-WIRED / CM-SEPOLICY-NEVERALLOW-WIRED + paired

§1.1 mutations (gate-code = separate work item).

Canonical authority: Constitution.md §11.4.110. Release blocker. No --
source-green-is-ready / --skip-clash-detector / --skip-coverage-
gate / --no-ready-verdict / --grep-proves-neverallow / --skip-
buildgraph-dryrun / --build-without-ready-verdict escape.

2026 06 03

§ . . – Resolve-by-stable-name-not-by-enumeration-index mandate (research-derived, – –)

Forensic anchor (genericised, 2026-06-03): a platform bound an audio output to a kernel-enumerated index (`card=0`); a second device of a different class enumerated FIRST at boot + took slot 0, shifting the intended device to slot 1 — the static index binding now pointed at the wrong device (policy mis-attached the sink, mis-assigned its TYPE, the output switcher labelled the AV receiver “Wired Headphone” + routing collapsed to stereo). The lower layer of the SAME stack already resolved by **name** (controller-name registry) — proving the brittleness was the *index*, not the resolution capability. Generalises to every enumerated resource whose ordinal is assigned at discovery/boot/hotplug (non-deterministic across reboots, additions, topology changes). The mandate: any binding to a hardware device / resource handle / enumerated entity (audio cards, display connectors, network interfaces, storage devices, GPU render nodes, input/camera devices, container/process slots) MUST resolve by a **stable identifier** (name / UUID / serial / label / controller-name / content-hash / sink-reported identity), NOT by enumeration index / ordinal / slot, UNLESS the platform documents the ordinal as deterministically pinned AND the pin is captured + asserted in the binding. Where a stable identifier exists at one layer, every layer binding the same resource MUST use it (mixed by-name-here / by-index-there is the forbidden weak link). Honest boundary (§11.4.6): when only an ordinal exists, pin deterministically via the platform’s mechanism, capture the pin in the §11.4.108 runtime signature, document residual fragility as UNCONFIRMED: -class — never silently trust an unpinned ordinal. Classification: universal (§11.4.17) — platform-neutral binding-robustness discipline; the project supplies its stable-identifier mechanism (ALSA card name, DRM connector name, NIC predictable name, block-device UUID, etc.)

+ the layers that must agree per §11.4.35. Composes §11.4.6/.8/.69/.108/.110. Propagation gate CM-COVENANT-114-111-PROPAGATION (literal 11.4.111) + recommended CM-RESOLVE-BY-NAME-NOT-INDEX + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: Constitution.md §11.4.111. Release blocker. No --allow-index-binding / --ordinal-is-stable-enough / --skip-name-resolution / --trust-unpinned-index escape.
11 4 112
2026 06 03

§ . . — Structural-impossibility won't-fix classification mandate (research-derived, — —)

Forensic anchor (genericised, 2026-06-03): deep research (§11.4.8) into relocating protected (content-protection / secure-surface) video to a secondary display PROVED — from authoritative platform/HDCP docs + reproducible captured behaviour (a secure surface is blanked on any output lacking the secure flag; mirror/screenshot of a secure layer returns black) — that the goal is **structurally impossible by platform design**, not a missing feature. Without a durable classification it is re-investigated every cycle (re-read the same sources, re-run the same probes, re-derive the same impossibility — compounding wasted effort). The mandate: when deep research per §11.4.8 PROVES (cited authoritative sources AND, where applicable, reproducible captured evidence) a goal is structurally impossible on the target platform (forbidden by platform design / hardware-protocol constraint / documented kernel-or-API limitation — NOT merely unimplemented or hard), the goal MUST be: (1) classified Won't-fix + closed per §11.4.90 with reason structurally-impossible ; (2) documented with the impossibility evidence — cited source URLs (per §11.4.99) + reproducible probe/captured evidence — in the tracker entry + relevant docs/ guide; (3) NOT re-attempted in future cycles — a reopen MUST cite NEW evidence the constraint changed (per §11.4.34 + §11.4.7), never re-derive the same impossibility; (4) paired with the correct posture — the entry states what the project DOES instead, so “impossible” is never confused with “broken/unhandled”. Honest boundary (§11.4.6): reserved for *proven* impossibility — “could not find a way” / “very hard” / “no time” are Operator-blocked (§11.4.21) or open work, NOT won't-fix; mislabelling them is a §11.4 planning-layer bluff. A future platform change can make the impossible possible (durable, not eternal). Classification: universal (§11.4.17) — platform-

neutral effort-conservation + honesty discipline; the project supplies the impossible goal, its constraint + cited evidence per §11.4.35. Composes §11.4.6/.7/.8/.34/.90/.99. Propagation gate `CM-COVENANT-114-112-PROPAGATION` (literal `11.4.112`) + recommended `CM-WONT-FIX-STRUCTURAL-IMPOSSIBILITY` + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: `Constitution.md` §11.4.112. Release blocker. No
`--wont-fix-without-proof` / `--reattempt-closed-impossible` / `--skip-impossibility-evidence` / `--impossible-equals-broken` escape.
11 4 113
2026 06 03

§ . . — Absolute no-force-push + merge-onto-latest-main mandate (User mandate, — —)

Forensic anchor — verbatim user mandate (2026-06-03): “Any force-push is strictly forbidden! We must for every Submodule take as a base latest commit on Submodule’s main (or master) branch, then on top of it carefully to merge all changes that have to be pushed! Once all merging is carefully done we perform commit and push to all Submodule’s upstreams!” Force-push is STRICTLY FORBIDDEN with NO exception — `git push --force`, `--force-with-lease`, `+<ref>`, any history-rewriting overwrite of a remote ref — against EVERY repository this Constitution governs (main repo, this constitution submodule, every owned + nested submodule, every upstream). No operator-approval path, no “after merge-first audit” path. Mandated 6-step integration for any repo/submodule with commits to publish OR diverged mirrors: (1) `git fetch --all --prune --tags` all remotes; (2) base = LATEST commit on canonical `main` / `master` (most-advanced mirror tip); (3) carefully MERGE every change on top — union, preserve BOTH sides, NEVER `-s ours` / rebase / reset that drops commits (§9 no-commit-loss); (4) resolve conflicts carefully — no markers, no file dropped, gates/tests still pass; (5) commit the merge (stage only intended files, NEVER `git add -A` in a submodule per §11.4.30); (6) push to ALL upstreams — each is a fast-forward because the merge commit descends from every mirror tip, so NO force is needed (an upstream still rejecting → return to step 1 for it, merge its new tip, re-validate, re-push). TIGHTENS §11.4.41 / §11.4.71 / §9.2 / CONST-043 — force-push escape hatch REMOVED: even WITH operator approval, even after a clean merge-first audit, force-push is forbidden, because the merge-onto-latest-main path is always available so force is never necessary; those clauses’ merge-first/fetch-first

machinery stays, only their terminal “...then force-push” step is struck.

Classification: universal (§11.4.17). Composes §2.1 / §9 / §9.2 / §11.4.4 / §11.4.6 / §11.4.26 / §11.4.37 / §11.4.40 / §11.4.41 / §11.4.71 / §11.4.88 / CONST-043.

Propagation gate CM-COVENANT-114-113-PROPAGATION (literal 11.4.113) + recommended CM-NO-FORCE-PUSH-ABSOLUTE (scan tracked scripts/hooks for push --force / --force-with-lease / push +<ref> + reject; §11.4.109-class PreToolUse guard blocks the class) + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: Constitution.md §11.4.113. Release blocker. No --force / --force-with-lease / --allow-force-push / --force-push-
authorised / --skip-merge-onto-main escape.

§ . . — Last-known-good-tag
regression isolation mandate (. . -dev
remediation, - -)

When a previously-working feature/behaviour is observed broken, the FIRST diagnostic action MUST be to identify the last release tag (or commit/build/deploy) at which it was KNOWN-GOOD and diff/bisect the broken state against it — BEFORE any open-ended root-cause hunt or speculative fix. The known-good revision is the regression oracle: bounds the search to commits between good and now (git diff <good-tag>..HEAD --stat of the feature's files = captured evidence), marks it a regression REPAIR not from-scratch design, gives each suspect file a behavioural oracle. An operator-volunteered known-good tag is load-bearing → act on it first. Default to SURGICAL forward-fix (keep post-good-tag features, revert ONLY the broken sub-part) over wholesale revert (which loses the batch's other working features) unless operator prefers wholesale. Honest boundary (§11.4.6): “it worked before” is a HYPOTHESIS until the tag is identified AND the feature confirmed working there; “probably regressed last batch” without the diff is a guess; a never-worked feature is not a regression. Classification: universal (§11.4.17). Composes §11.4.4 / §11.4.6 / §11.4.7 / §11.4.40 / §11.4.43 / §11.4.102 / §11.4.108. Propagation gate CM-COVENANT-114-114-PROPAGATION (literal 11.4.114) + recommended CM-REGRESSION-ISOLATED-AGAINST-KNOWN-GOOD + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: Constitution.md §11.4.114. Release blocker. No
1 1 4 1 1 5
--skip-known-good-diff / --root-cause-from-scratch / --assume-
regression-source / --wholesale-revert-without-isolation 1 1 8
2 0 2 6 0 6 0 3 escape.

§ . . – RED-baseline-on-the-broken-artifact + polarity-switch mandate (. . -dev remediation, - -)

Strict refinement of §11.4.43's RED step. Every RED test MUST REPRODUCE the defect on the CURRENT pre-fix artifact (the actual broken build/deploy), capturing positive evidence per §11.4.5 / §11.4.69 / §11.4.107 that the defect is genuinely present — never a synthetic failure the fix is then written to agree with. The SAME source MUST carry a single polarity switch (env flag / parameter, canonical RED_MODE , default 1 = reproduce-and-assert-defect-present) that flipped to 0 post-fix becomes the GREEN regression-guard asserting the defect is ABSENT. One source, two roles: the bug-catcher IS the regression-guard; no separate happy-path test is the primary guard (that only proves the test agrees with the fix — the §11.4.43 PASS-bluff). RED-on-broken then GREEN-on-fixed (clean target per §11.4.108) both captured. Honest boundary (§11.4.6): a RED run that does NOT fail on the broken artifact is a finding (close per §11.4.7 with negative evidence, or fix the blind test). Classification: universal (§11.4.17). Composes §11.4.1 / §11.4.2 / §11.4.5 / §11.4.69 / §11.4.107 / §11.4.4 / §11.4.7 / §11.4.43 / §11.4.50 / §11.4.108 / §11.4.114. Propagation gate CM-COVENANT-114-115-PROPAGATION (literal 11.4.115) + recommended CM-RED-POLARITY-SWITCH-PRESENT + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: Constitution.md §11.4.115. Release blocker. No --
1 1 4 1 1 6
green-only-test / --skip-red-reproduction / --separate-happy-path-
suffices / --synthetic-red-OK 1 1 8
2 0 2 6 0 6 0 3 escape.

§ . . – Real-time conductor autonomous-test-framework sync channel mandate (. . -dev remediation, - -)

Any autonomous, long-running test/QA/validation framework an external orchestrator (conductor agent / operator) depends on for real-time decisions MUST expose a real-time sync channel: (1) a structured append-only event stream

(JSONL or equivalent, one event per line, never rewritten — session-start / phase-transition / per-test-or-challenge-start / captured-evidence-path / external-call (LLM/vision/sink-probe) / error / per-item-verdict); (2) an atomically-rewritten status snapshot (single small file, write-temp-then-rename so no torn read) with current session/phase/item + counters + last verdict. Verdicts use the closed PASS / FAIL / SKIP / OPERATOR-BLOCKED vocabulary (§11.4.45). The conductor tails it live → real-time sync, §11.4.4 interrupt on a fresh defect, never idles blindly (§11.4.94 / §11.4.97). Anti-bluff: every verdict event MUST carry the evidence path that backs it (§11.4.69) — a PASS event with no evidence path is a channel-layer bluff; a snapshot PASS contradicted by a stream with no evidence event → treat as FAIL; no start-event → no verdict-event. Owned project-agnostic submodule (§11.4.28): channel stays project-neutral, the consumer registers its data via the public API at runtime, never hardcoded. Classification: universal (§11.4.17). Composes §11.4.4 / §11.4.5 / §11.4.69 / §11.4.27 / §11.4.28 / §11.4.45 / §11.4.52 / §11.4.89 / §11.4.94 / §11.4.97. Propagation gate `CM-COVENANT-114-116-PROPAGATION` (literal `11.4.116`) + recommended `CM-AUTONOMOUS-FRAMEWORK-SYNC-CHANNEL` + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: `Constitution.md` §11.4.116. Release blocker. No `--no-sync-channel` / `--end-of-session-report-suffices` / `--verdict-without-evidence-path` / `--torn-status-write-OK` escape.

§ . . — Computer-vision / OCR pixel-oracle fallback for non-introspectable UIs
mandate (. . -dev remediation,
— —)

Any test that drives a UI control OR asserts on-screen content MUST NOT assume the accessibility/semantic/DOM hierarchy is the source of truth for what the user sees. When the hierarchy is blank/partial/known-unreliable (TV-Compose, leanback, canvas/ SurfaceView /GL, games, custom-rendered UIs), the test MUST fall back to a PIXEL ORACLE: (1) DRIVE input by computer-vision template-match (locate a control by rendered appearance → tap coordinates), not a node id; (2) ASSERT content by ROI OCR (read the rendered text the user reads) with a per-word confidence floor + ROI per §11.4.107(12), not a hierarchy text attribute that may be absent/stale. The tool MUST both drive input AND read pixels — a hierarchy-only tool is NOT a content oracle. Makes §11.4.52's "near-empty

hierarchy → INFEASIBLE” constructive: pixel-drive, not only SKIP. Anti-bluff (§11.4.107(10)): the CV/OCR analyzer is self-validated (golden-good PASS, golden-bad FAIL, wired into meta-test); thresholds calibrated on the project’s own frames, not hardcoded (§11.4.6). Honest boundary: BOTH hierarchy blank AND pixel oracle infeasible (secure surface black-captures §11.4.112, geo-unreachable) → SKIP-with-reason §11.4.3 + tracked operator-attended migration §11.4.52, never fake PASS. Classification: universal (§11.4.17). Composes §11.4.5 / §11.4.6 / §11.4.48 / §11.4.49 / §11.4.52 / §11.4.107 / §11.4.112. Propagation gate CM-COVENANT-114-117-PROPAGATION (literal 11.4.117) + recommended CM-CV-OCR-PIXEL-ORACLE-FALLBACK + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: Constitution.md §11.4.117. Release blocker. No `--
11 4 118
hierarchy-is-content-oracle / --skip-pixel-fallback / --unvalidated-
ocr-OK / --hardcoded-ocr-threshold-OK escape.
1 1 8 2026 06 03`

§ . . – Discovery-pressure to confirm known-issue-set completeness mandate (. . -dev remediation, – –)

A remediation/release cycle MUST NOT treat “every reported defect is fixed” as “the build is good” — the reported set is biased by what the operator happened to test (“we only see what we test”). After/alongside fixing the reported set, run a discovery + stress pass across ALL target devices/environments that deliberately exercises subsystems/journeys/edge-cases BEYOND the reported defects — to CONFIRM the reported set is the COMPLETE critical set + surface unreported defects before the user does. The pass MUST produce PROVABLE coverage: an enumerated list of subsystems/journeys/stress-scenarios exercised, each with outcome (no-new-issue, or a new tracker entry per §11.4.15 / §11.4.16). “We found no other issues” is a §11.4 bluff unless paired with “here is the enumerated set we exercised” — absence of evidence of looking is not evidence of absence. New findings trigger §11.4.4 interrupt + §11.4.114/§11.4.115 isolation → RED → fix. Honest boundary (§11.4.6): discovery reduces the unknown-unknown surface, does not prove zero remaining defects — earned claim is “reported set + enumerated discovery set addressed,” un-exercised subsystems stated as honest coverage gaps (§11.4.3), never silently implied clean. Classification: universal (§11.4.17). Composes §11.4.4 / §11.4.5 / §11.4.69 / §11.4.25 / §11.4.40 / §11.4.42 / §11.4.52 /

§11.4.85 / §11.4.114 / §11.4.115 / §11.4.119. Propagation gate CM-COVENANT-114-118-PROPAGATION (literal 11.4.118) + recommended CM-DISCOVERY-COVERAGE-ENUMERATED + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: Constitution.md §11.4.118. Release blocker. No --reported-set-suffices / --skip-discovery-pass / --no-issues-without-coverage-list / --assume-complete escape.

2026 06 03

§ . . – Single-resource-owner
partitioning for parallel hardware testing
mandate (. . -dev remediation,
– –)

Strict refinement of §11.4.58 / §11.4.103 for hardware contention. When multiple parallel streams exercise SHARED hardware / any exclusive-access resource (a media-playback path, a single HDMI/audio sink, an exclusive device handle, a serial/JTAG line, a GPU under exclusive capture), exactly ONE stream MUST own each such resource at a time. The owner drives it (playback, input injection, capture); every other concurrent stream targeting it MUST be READ-ONLY (passive probes — dumpsys / /proc / /sys reads / sink-side network probes / log tails). Parallelism partitioned by resource: distinct devices/sinks/handles fully concurrent (stream-per-device), but the same device's exclusive resource is single-owner. Ownership enforced by an advisory lock/token (§11.4.58 L3 + the hardware analogue of §11.4.84 quiescence), event-driven (claim when the resource frees, release the moment work completes so the next queued stream claims it per §11.4.103 auto-backfill). Why: concurrent drivers of one exclusive resource produce CROSS-CONTAMINATED evidence (sink reports the wrong stream's audio, foreground belongs to whichever am start landed last, input events interleave) — a PASS under contention is a §11.4 evidence-integrity bluff. Honest boundary (§11.4.6): a “read-only” probe stream MUST issue NO state-changing command against a device it does not own; a non-partitionable resource → streams serialize (single-owner over time), not run concurrently and hope. Classification: universal (§11.4.17). Composes §11.4.5 / §11.4.69 / §11.4.13 / §11.4.50 / §11.4.58 / §11.4.82 / §11.4.84 / §11.4.103 / §11.4.118. Propagation gate CM-

COVENANT-114-119-PROPAGATION (literal 11.4.119) + recommended
CM-SINGLE-RESOURCE-OWNER-PARTITION + paired §1.1 mutation (gate-code =
separate work item).

Canonical authority: Constitution.md §11.4.119. Release blocker. No --
allow-concurrent-resource-drivers / --skip-ownership-lock / --read-
only-may-mutate / --contended-evidence-OK escape.
11 4 120
2026 06 03

§ . . – Fix-breaks-its-own-gate
reconciliation mandate (. . -dev
remediation, – –)

When a correct fix causes a pre-existing gate/test to FAIL because that gate asserted the OLD (now-removed-or-changed) behaviour, the gate FAIL is the CORRECT signal the fix landed — it MUST NOT be suppressed by either forbidden response: (1) FAKE-PASSING the gate (edit to always pass , weaken to a tautology, or delete — a §11.4 gate-layer bluff + a §11.4.84 mutation-residue risk); (2) REVERTING the correct fix to satisfy the stale gate (re-introducing the defect). Required response = RECONCILIATION: rewrite the gate to assert the NEW mechanism the fix introduced, backed by captured evidence of the new correct behaviour, AND update its paired §1.1 mutation so the mutation breaks the NEW invariant. Discriminator vs bluffing: after reconciliation gate + mutation still form a valid §1.1 pair (mutate the new invariant → gate FAILs); a bluffed gate's mutation no longer makes it FAIL. The reconciliation MUST be a visible evidence-cited change, never a silent assertion-weakening. Honest boundary (§11.4.6): a post-fix gate FAIL is NOT automatically “stale gate” — investigate per §11.4.102 first; the FAIL may be the gate correctly catching a REGRESSION the fix introduced (then the FIX is wrong). Reconcile ONLY when investigation PROVES the gate asserted old-correct-now-removed behaviour AND the new behaviour is intended + evidence-confirmed. Classification: universal (§11.4.17). Composes §11.4.1 / §11.4.4 / §11.4.6 / §11.4.84 / §11.4.102 / §11.4.108 / §1.1. Propagation gate CM-COVENANT-114-120-PROPAGATION (literal 11.4.120) + recommended CM-GATE-RECONCILED-NOT-FAKE-PASSED + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: Constitution.md §11.4.120. Release blocker. No

`--fake-pass-stale-gate / --revert-fix-for-gate / --weaken-assertion-
to-pass / --delete-failing-gate escape.`

§ . . — No-commit-while-build-writes-tracked-artifacts mandate (. . -dev remediation, - -)

A commit (especially `git add -A` / any broad stage) MUST NOT run while a build/packaging/generation step is actively writing artifacts INTO tracked (version-controlled) directories — it races the writer and stages a PARTIAL or stale artifact (a §11.4.108 SOURCE → ARTIFACT integrity failure landed in version control). The commit MUST be deferred until the build step that writes tracked artifacts has COMPLETED, so the tree is quiescent at the artifact layer AND the committed artifacts are the FRESH, whole outputs (no stale pre-rebuild artifact, no half-written file). Before committing tracked build outputs, verify the writing step finished (process exit / completion marker / per-artifact mtime ≥ build-start). Build-output analogue of §11.4.84 (no commit while a mutation gate is in flight): both close “commit captures transient non-final tree state” at two write-sources. Gitignored build outputs (`out/` / `dist/`) are exempt — the race doesn’t apply. Honest boundary (§11.4.6): “the build probably finished” is not “the build finished” — verify with a completion signal; source changes NOT touching the build’s tracked-artifact dirs are fine mid-build. PROJECT-SPECIFIC instance (which tracked dir + which build steps write it) recorded in the consuming project’s own governance per §11.4.35. Classification: universal (§11.4.17). Composes §11.4.6 / §11.4.30 / §11.4.58 / §11.4.84 / §11.4.88 / §11.4.96 / §11.4.103 / §11.4.108. Propagation gate `CM-COVENANT-114-121-PROPAGATION` (literal `11.4.121`) + recommended `CM-NO-COMMIT-DURING-ARTIFACT-WRITE` + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: Constitution.md §11.4.121. Release blocker. No `--commit-during-build / --stage-partial-artifact-OK / --assume-build-finished / --skip-build-completion-check escape.`

§ . . — No-silent-removal-of-existing-components-without-operator-confirmation mandate (User mandate, — —)

Forensic anchor — verbatim user mandate (2026-06-03): “Never ever remove any application, system component or service from already existing codebase / System without interactively asked question to us! THIS IS MANDATORY RULE / CONSTRAINT!” Case study (FACT): in the 1.1.8-dev burn-down, F2 (an Apple-TV-class app) + F4 (a Huawei HMS component) were removed WITHOUT asking; the operator reversed both.

No application / system component / service / package / feature / driver / module / library / prebuilt asset — any already-existing end-user capability — may be removed (deleted, dropped from the package set, disabled-into-non-shipping, un-bundled, de-listed) from the existing codebase / shipped System WITHOUT FIRST interactively asking the operator (via the §11.4.66 interactive mechanism —

`AskUserQuestion` on Claude Code, never free-text, never silent, never an autonomous decision) and receiving an EXPLICIT keep-or-remove decision. A removal is: a deletion from the package/service/module manifest set whose NET EFFECT is “a capability that shipped before no longer ships.” Adding / replacing-with-operator-approved-equivalent / fixing is NOT a removal; when uncertain, treat AS a removal and ask (§11.4.6 no-guessing). A silent removal is a release blocker regardless of rationale (dedup / “broken anyway” / geo-restricted / incompatible / superseded). The tracked DROP path: ask → operator approves → mark item `Obsolete (→ Fixed.md)` with reason `feature-removed` + operator-approval citation (§11.4.90) → then remove. Classification: universal (§11.4.17). Composes §11.4.66 / §11.4.101 (removal of a live capability is exactly the high-blast-radius class §11.4.101 reserves for an operator question, never an autonomous decision) / §11.4.90 / §11.4.112 / §11.4.6 / §11.4.40 / §11.4.42. Propagation gate `CM-COVENANT-114-122-PROPAGATION` (literal `11.4.122`) + recommended gate `CM-NO-SILENT-COMPONENT-REMOVAL` + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: `Constitution.md` §11.4.122. Release blocker. No `--remove-without-asking` / `--silent-removal` / `--autonomous-removal-OK` / `--dedup-removal-exempt` / `--it-was-broken-anyway` escape.

§ . . — Rock-solid-proof-or-deep-research mandate (User mandate, — —)

Forensic anchor — verbatim user mandate (2026-06-03): “Every single reported issue MUST BE fully and 100% validated with rock solid proofs! Nothing can be considered fixed or completed without hard evidence! No false results or bluff(s) of any kind is allowed! If we are not sure on how to achieve full testing, validation and verification of something we MUST ALWAYS perform deep web research for all possible data (articles, documentation, guides, and other resources) and opensourced codebases which we can use to solve our problems and perform testing with validation and verification which produces rock-solid evidence(s) and leaves no space for false results or any kind of bluff!” Case study (FACT): in the 1.1.8-dev remediation the validation method for FLAG_SECURE secondary-display capture (black frames) + non-introspectable streaming-app UIs (blank accessibility tree) was unclear; deep web research (docs/research/testing_frameworks_20260603/) yielded the CV/OCR/liveness/sink-probe oracle stack (now §11.4.107/§11.4.112/§11.4.117), making rock-solid evidence possible — “unclear how to validate” is a research trigger, never a bluff licence.

Every reported issue / fix / claimed completion MUST be 100% validated with rock-solid CAPTURED proof (§11.4.5/§11.4.69/§11.4.107) before it is marked fixed/implemented/completed (§11.4.33). Nothing is fixed/complete without hard captured evidence — metadata-only / config-only / absence-of-error / grep-without-runtime PASS forbidden (§11.4/§11.4.1); no false results, no bluff of any kind. Operative addition: when UNSURE how to fully test/validate/verify something (no obvious evidence-producing method, OR the candidate yields only metadata/config/absence-of-error evidence), the agent MUST ALWAYS FIRST perform deep web research per §11.4.8+§11.4.99 (official docs, articles, guides, vendor refs, standards, issue trackers, reusable open-source codebases) to DISCOVER or BUILD an evidence-producing validation method that leaves no room for a false result. Declaring “untestable”/“not automatable” or accepting a metadata-only PASS WITHOUT first exhausting this deep-research path is itself a §11.4.123 violation (same severity as a PASS-bluff); the cited research (URLs + the method, OR “NO external solution found — original work” per §11.4.8) is the proof the path was exhausted. Only after that research genuinely fails may the item be

PENDING_FORENSICS: / operator-blocked (§11.4.21)/ structurally-impossible won't-fix (§11.4.112) with the cited research as earned evidence.

Classification: universal (§11.4.17). Composes §11.4.5/§11.4.6/§11.4.8/§11.4.52/§11.4.69/§11.4.99/§11.4.107/§11.4.118/§11.4.21/§11.4.112. Propagation gate CM-COVENANT-114-123-PROPAGATION (literal 11.4.123) + recommended gate CM-ROCK-SOLID-PROOF-OR-RESEARCH + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: Constitution.md §11.4.123. Release blocker. No --metadata-pass-suffices / --skip-proof /
11 4 124
--untestable-without-research / --config-only-closure-OK / --bluff-when-unsure escape.
2026 06 04

§ . . — Dead/unwired-code investigate-before-remove mandate (User mandate, — —)

Forensic anchor — verbatim user mandate (2026-06-04): “Before removing any seemingly-dead (zero-importer / unwired) codebase, we MUST investigate via git history where/how it was originally used and how it became dead. Removal is permitted ONLY when we have captured PROOF it is genuinely no longer needed — and that removal MUST be its own separate commit with a proper descriptive message. If there is no such proof, the code MUST be investigated for where/how it should be wired in properly, and any missing or unwired tests MUST be added. We MUST ALWAYS be extra careful with any codebase removal.”

“Zero importers / never called / unwired ⇒ dead ⇒ delete” is a GUESS (§11.4.6), never a finding — a “no references” result proves only current non-reference, not genuinely-unneeded. Before removing ANY seemingly-dead element (zero-importer / never-called / unwired function / method / type / file / module / package / asset / config / build target) the agent MUST FIRST investigate via git history (git log --follow , git log -S / -G pickaxe, blame on the deleted call-site) and capture as FACT: (1) WHERE/HOW it was originally wired in, (2) WHEN/HOW it became dead (call-site deleted deliberately / by mistake-regression / never-completed / refactored-unreachable), (3) whether “no references” is real OR a hidden reference the static tool cannot see (reflection / dynamic dispatch / build-tags / codegen / DI / plugin registry / FFI / config-driven). Removal is permitted ONLY with captured PROOF the element is genuinely no longer needed; that removal MUST be its OWN SEPARATE COMMIT (independently reviewable + revertible per §11.4.84/§11.4.92) with a descriptive message citing the git-history

evidence (+ §11.4.122 operator-confirmation if it is an end-user capability; tracked via §11.4.90 `obsolete` (→ `Fixed.md`)). With NO such proof the element MUST NOT be removed — instead investigate WHERE/HOW to wire it in properly (restore a mistakenly-deleted call-site per §11.4.114; finish never-completed wiring) AND add any missing / unwired tests (§11.4.27/§11.4.43/§11.4.115). ALWAYS be extra careful with removal: when uncertain, default to NOT removing (investigate + wire + test) per §11.4.6 + §11.4.101 + §11.4.122; “probably dead” is never sufficient — the bar is captured proof. Classification: universal (§11.4.17). Composes §11.4.6 / §11.4.8 / §11.4.84 / §11.4.90 / §11.4.92 / §11.4.101 / §11.4.114 / §11.4.122 / §11.4.27 / §11.4.43 / §11.4.115. Propagation gate `CM-COVENANT-114-124-PROPAGATION` (literal `11.4.124`) + recommended gate `CM-DEAD-CODE-INVESTIGATE-BEFORE-REMOVE` (a net-deletion commit must be removal-only + cite the git-history investigation OR be part of a tracked `Obsolete` item) + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: `Constitution.md` §11.4.124. Release blocker. No

```
--zero-importers-means-dead / --delete-unwired-on-sight /
1 1 4 1 2 5
--skip-git-history-investigation / --remove-without-proof /
--bundle-removal-with-other-work escape.
```

2 0 2 6 0 6 0 4

§ . . — Code-review-agent gate
before pre-build + main build (mandatory multi-layer review) (User mandate,
- -)

Forensic anchor — verbatim user mandate (2026-06-04): “After all fixes/changes/implementations are done, BEFORE running pre-build tests and the main build, dispatch code-review agent(s) that analyze all work done + all existing data/facts + the existing codebase + current git history to determine quality, safety, and whether the fixes/changes will REALLY work; they MUST validate and verify that every test covering the fixes/changes genuinely validates the work with NO chance of false results or bluff of any kind. Any finding MUST be fixed, polished, improved, and covered with additional tests before the build proceeds. Multiple strong layers of checks.”

After all fixes / changes / implementations in a batch are done, and BEFORE running the pre-build test sweep AND the main (artifact) build (for ANY project), the agent MUST dispatch one or more dedicated code-review agent(s) (subagent-

driven by default per §11.4.70/§11.4.20) performing a multi-layer review that: (1) analyzes ALL work done in the batch (every fix/change + its source diff + stated intent); (2) analyzes ALL existing data + facts (captured evidence per §11.4.5/ §11.4.69/§11.4.107, tracker entries, prior findings, the §11.4.108 runtime-signature registry); (3) analyzes the existing codebase (blast radius per §11.4.92, cross-feature interaction, contract integrity of every dependency); (4) analyzes current git history (what each change touched, how it composes with concurrent/recent work, whether it reproduces a known-broken pattern per §11.4.114/§11.4.124); (5) determines quality + safety + will-it-REALLY-work (robust + not error-prone — no solve-A-create-B; no host/data/security regression; genuinely delivers the end-user-visible behaviour per §11.4/§107); (6) validates + verifies the tests covering the work — every covering test genuinely exercises the work-under-test and catches its negation, with ZERO chance of a false result or bluff (a test that PASSES on broken-for-the-user work, a metadata-only/config-only/absence-of-error/grep-without-runtime assertion, or a gate whose paired §1.1 mutation does not make it FAIL is a finding). Any finding (defect / error-prone change / safety risk / will-not-really-work / bluff-or-false-result-capable test / missing-coverage gap) MUST be fixed, polished, improved, and covered with additional tests (four-layer per §11.4.4(b), TDD-RED-first per §11.4.43/§11.4.115) BEFORE the pre-build sweep + main build proceed; the review iterates (re-review after each remediation) until no blocking findings remain. The review is itself anti-bluff (its conclusions are captured evidence per §11.4.5/§11.4.69; a rubber-stamp review of a defective batch = PASS-bluff). It is one of MULTIPLE STRONG LAYERS — complementing, never replacing, the §1 pre-build sweep, §11.4.92 multi-pass (author-side self-review; §11.4.125 adds the structurally-separated reviewer seam per §11.4.70), §11.4.108 four-layer fix-verification, §11.4.110 build-readiness verdict, and the post-build / runtime-on-clean-target / user-visible layers. Composes §11.4 / §11.4.1 / §11.4.4 / §11.4.6 / §11.4.40 / §11.4.43 / §11.4.50 / §11.4.70 / §11.4.20 / §11.4.92 / §11.4.102 / §11.4.107 / §11.4.108 / §11.4.110. Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-125-PROPAGATION` (literal `11.4.125`) + recommended gate `CM-CODE-REVIEW-GATE-BEFORE-BUILD` (build starts only with a fresh code-review-completed marker for the current batch, produced after the last fix + before the pre-build sweep + main build) + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: Constitution.md §11.4.125. Release blocker. No
1 1 4 1 2 6
--skip-code-review / --build-without-review / --no-review-gate / --
review-optional / --trust-the-author escape.
2 0 2 6 0 6 0 4

§ . . — Default autonomous-loop
working mode from first prompt (User mandate,
— —)

Forensic anchor — verbatim user mandate (2026-06-04): “Make sure that you continue work in endless fully autonomous loop, do not stop until new fully validated and verified version (tag) is created and published (all submodules and main repo) or IN A CASE OF some other main stream work until it is fully completed with all side work streams and nothing else is left in our working queue! THIS MUST BE ALWAYS the default working mode without us asking you! We tend to achieve ABSOLUTE EFFICIENCY, with this and all other projects which will incorporate this MANDATORY RULE / CONSTRAINT!!! This way of (your) working will be ALWAYS applied / followed / executed / fully respected, as soon as we assign / send first request (prompt) in the session! This stops only if we explicitly say so or nothing is left to be done in current working scope (release that will come / upcoming version)!!! Any mimicking (imitation) of this behavior / rules / mandatory constraints, false results or any kind of bluff(s) is ABSOLUTELY FORBIDDEN!!!”

The endless fully-autonomous loop is the **DEFAULT working mode**, engaged automatically the moment the operator sends the FIRST request / prompt of a session — the operator MUST NOT have to ask for it, request it, restate it, or re-enable it per session. §11.4.87 framed the endless-loop covenant as an explicit-instruction opt-in; §11.4.126 is the **capstone** that promotes the same covenant to always-on: from the first prompt onward, every agent operates in the §11.4.87 loop discipline as the standing default, with §11.4.94 zero-idle, §11.4.97 maximum-idle-use, §11.4.101 autonomous-decision-over-blocking, and §11.4.103 continuous-parallel-stream all engaged by default — no per-session activation handshake. The continuation contract: the loop continues until ONE of two terminal conditions holds — (A) **Release scope** — a new, fully-validated-and-verified version (tag) is created AND published across all owned submodules AND the main repo to all configured remotes (per §2.1 + §11.4.40 + §11.4.113); OR (B) **Non-release main-stream scope** — the main-stream goal is fully completed AND every side work stream is done AND the working queue holds nothing left for the current scope. Until (A) or

(B) holds, the agent MUST keep working per §11.4.42 / §11.4.72 / §11.4.94 / §11.4.103. The loop STOPS ONLY on: (1) the operator explicitly saying so (STOP / pause / end); (2) nothing left to do in the current working scope with the queue genuinely empty per the (A)/(B) terminal conditions; (3) a §12 host-session-safety demand. Idle-while-blocked parks one work unit, it does not stop the loop (§11.4.101 + §11.4.94 + §11.4.97). Goal — ABSOLUTE EFFICIENCY; applies to this project AND every project that incorporates this Constitution. Anti-bluff: mimicking / imitating this loop behaviour, narrating continuation without performing it, fabricating progress, or emitting false / bluff results of ANY kind is ABSOLUTELY FORBIDDEN (composes §11.4 / §11.4.1 / §11.4.2 / §11.4.5 / §11.4.6 / §11.4.50 / §11.4.69 / §11.4.107); a report claiming the loop ran while no real work / no captured evidence was produced is a §11.4 PASS-bluff at the operating-mode layer. Classification: universal (§11.4.17). Composes with §11.4.87 / §11.4.94 / §11.4.97 / §11.4.101 / §11.4.103 / §11.4.66 / §11.4.6 / §11.4.40 / §11.4.42 / §11.4.72 / §11.4.113 / §2.1 / §12. Propagation gate CM-COVENANT-114-126-PROPAGATION (literal 11.4.126) + paired §1.1 mutation (gate-code = separate work item).

Canonical authority: Constitution.md §11.4.126. Release blocker. No --ask-before-continuing / --single-turn-only / --not-default-loop / --mimic-OK escape.

§ . . — Session-handoff resumption-prompt mandate (User mandate, — —)

Forensic anchor — verbatim user mandate (2026-06-06): “make sure that in situations like this now when new session is needed you ALWAYS prepera such sentence - which will be valid for particular moment and the phase of the project and enough for work to continue.”

When a fresh session is needed (context limits / degradation) OR the operator asks whether a new session is needed / requests a handoff, the agent MUST ALWAYS prepare + proactively provide a ready-to-paste **resumption prompt valid for that EXACT moment + project phase** — self-contained for ZERO-loss resume. Two variants: SHORT first-sentence (“Read <handoff docs> , then continue <goal> ...”) + FULL block on demand. MUST (1) point to live handoff docs (.remember/remember.md if present + docs/CONTINUATION.md §12.10, read FIRST + git fetch --all); (2) state PHASE + NEXT action + terminal goal; (3) embed exact

live-state anchors (build IDs/MD5, device/target serials, commit HEAD, in-flight PIDs+log paths, captured-evidence paths); (4) restate binding constraints (anti-bluff §11.4, no-force-push §11.4.113, exact version/naming, hardware/target gotchas); (5) be MOMENT-VALID never generic. Handoff docs current BEFORE the prompt (§12.10). Missing/stale/generic = §11.4.127 violation. Composes §12.10 / §11.4.6 / §11.4.66 / §11.4.87 / §11.4.103 / §11.4.126. Classification: universal (§11.4.17). Gate CM-COVENANT-114-127-PROPAGATION (literal 11.4.127) + paired §1.1 mutation.

Canonical authority: Constitution.md §11.4.127. Release blocker. No
--skip-handoff-prompt / --generic-prompt-OK / --no-resumption-sentence / --handoff-without-state escape.

§ . . — Always-on device-recording mandate (User mandate, — —)

Forensic anchor — direct user mandate (2026-06-06): ALWAYS live-record all available data from all test/manual-testing devices, EXTRA carefully (never harm device / performance / no side effects); raw NOT processed without need (token-conscious); raw ALWAYS git-ignored + code-intelligence-excluded; only curated evidence committed, only at release prep.

For EVERY test/debug device + every device under known manual testing, across EVERY transport (USB / wireless ADB / SSH / serial / network introspection API), ALWAYS live-record all analysable data: activities, all logs, perf metrics (CPU/mem/I/O/thermal/load), every sink-side report (§11.4.13), any live-changeable param. (1)

Side-effect-free — non-invasive read-only probes, bounded sampling + write-volume, observer-effect budget; a perturbing recorder = §11.4.128 violation. (2)

Background + parallel + subagent-driven (§11.4.103 + §11.4.70) — never blocks main stream. (3) **Token-conscious record-now-analyse-later** — raw NOT processed without need; only analyse-trigger = release-tag prep (§11.4.40/ §11.4.42) OR explicit operator ask. (4) **Raw git-ignored (w/ §11.4.77 regen declaration) + code-intelligence-excluded (§11.4.78/79)** — only curated evidence committed, only at release prep under docs/qa/<run-id>/ (§11.4.83).

(5) Layout

<recording-root>/YYYY-MM-DD/<combined main+submodules hash>/<DEVICE>_<SERIAL>/recording_NNN/<files> . (6) **Anti-bluff** — recorder-running-but-no-growing-corpus = §11.4 bluff; curated findings trace to real raw paths;

recorder health = captured evidence (§11.4.5/.69). Composes §11.4.2/.5/.13/.69/.40/.42/.70/.77/.78/.79/.83/.103/.119. Classification: universal (§11.4.17). Gate CM-COVENANT-114-128-PROPAGATION (literal 11.4.128) + rec.gate CM-DEVICE-RECORDING-ALWAYS-ON + paired §1.1 mutation.

Canonical authority: Constitution.md §11.4.128. Release blocker. No
--skip-recording / --record-without-layout / --commit-raw-corpus /
--index-raw-corpus / --analyse-corpus-always / --invasive-probe-OK
escape.

§ . . — Huge-blocker release protocol (User mandate, — —)

Forensic anchor — direct user mandate (2026-06-06): on a huge blocker during release validation: STOP all testing → fix ALL discovered issues → process all recorded data from last session → land rock-solid fixes → author NEW validation+verification tests of ALL supported test types → rebuild → reflash → RESTART full validation+verification of every fix/change from the last release tag to now — both devices in parallel, recorded, real physical proofs, no bluff.

On a HUGE BLOCKER (release-blocking-severity: core capability broken / regression invalidating the cycle / blast radius reaching the batch's other fixes), in order, NO shortcut: (1) STOP all testing every device (§11.4.4 STOP at release granularity). (2) Fix ALL discovered issues (not just the blocker) — root-cause §11.4.102 + isolate against last known-good tag §11.4.114. (3) Process all recorded data from last session (the §11.4.128(3) release-prep analyse-trigger). (4) Land rock-solid fixes §11.4.123 + §11.4.43/.115 + §11.4.9. (5) Author NEW tests of ALL supported types §11.4.27 + §11.4.85, anti-bluff + paired §1.1 mutation. (6) Rebuild (full, not module-only) + reflash CLEAN target §11.4.108. (7) RESTART full validation from last tag to now §11.4.40 — RESTART not resume — both/all devices parallel §11.4.103/.119, recorded §11.4.128, real physical proofs §11.4.5/.69/.107, no bluff. BINDS the existing release anchors for the huge-blocker case (STOP → fix-all → process-recordings → new-tests-all-types → rebuild → reflash → full-restart + restart-not-resume), citing not duplicating. Composes §11.4.4/.40/.42/.9/.27/.85/.102/.108/.114/.115/.123/.128/.103/.119.

Classification: universal (§11.4.17). Gate `CM-COVENANT-114-129-PROPAGATION` (literal `11.4.129`) + rec. gate `CM-HUGE-BLOCKER-FULL-RESTART` + paired §1.1 mutation.

Canonical authority: `Constitution.md` §11.4.129. Release blocker. No `--resume-after-blocker` / `--spot-validate-after-fix` / `--skip-recording-analysis` / `--skip-new-tests` / `--module-only-after-blocker` / `--single-device-restart` escape.

§ . . — Post-remediation validate-the-fix-FIRST-after-redeploy (User mandate, — —)

Forensic anchor — direct user mandate (2026-06-06): when a blocker found during release validation is fixed and a new artifact (rebuild / new flashing image / redeploy) is produced + the target reflashed, FIRST re-test the SPECIFIC last-failing features + validate the just-incorporated fixes BEFORE the broader / full validation.

When a blocker / critical failure found during release validation is FIXED and a new artifact is produced + the target reflashed / redistributed / updated, in order, NO shortcut: (1) re-test the SPECIFIC last-failing features FIRST (targeted guard tests for exactly the defects this fix addressed) BEFORE any broader / full-suite validation; (2) validate the just-incorporated fixes with real captured evidence — the §11.4.115 RED test flips GREEN at `RED_MODE=0` on the new artifact AND the §11.4.108 runtime-signature verifies on the CLEAN target the redeploy produced (metadata-only/config-only/absence-of-error/grep-without-runtime PASS forbidden §11.4/.1; proof §11.4.5/.69/.107/.123); (3) ONLY after the targeted fix is CONFIRMED working proceed to the §11.4.40 full retest from last tag to now. Rationale: a first fix attempt may not work / be incomplete / regress again under the new artifact — confirming FIRST catches fix-did-not-take immediately instead of hours later at the END of a full cycle (then restarting §11.4.129); cheap-confirmation-first is §11.4.82 applied to the post-blocker reflash. This is §11.4.46 specialised for the post-blocker-reflash case + the targeted-confirmation phase that GATES §11.4.129's step-7 full-restart. Honest boundary §11.4.6: “probably took” ≠ “took”; a still-FAILing targeted re-test re-enters the §11.4.114/.115 isolate → RED → fix loop, never proceeds to the full cycle on a still-broken fix.

Composes §11.4.4/.40/.46/.108/.114/.115/.123/.129/.82. Classification: universal (§11.4.17). Gate CM-COVENANT-114-130-PROPAGATION (literal 11.4.130) + rec. gate CM-FIX-FIRST-AFTER-REDEPLOY + paired §1.1 mutation.

Canonical authority: Constitution.md §11.4.130. Release blocker. No

```
11 4 131
--skip-targeted-retest / --full-cycle-first / --assume-fix-took / --
validate-fix-at-end / --skip-red-green-flip-on-new-artifact escape.
2026 06 07
```

§ . . — Standing session-resumption
file mandate (User mandate,
— —)

Forensic anchor — verbatim user mandate (2026-06-07): “Make this markdown a standard file which will be written EVERY TIME when we need fresh session out of the box! It MUST BE always up to date and in sync so whenever new session is created all we have to do is just point to it!”

Every project MUST maintain a SINGLE canonical, always-current **session-resumption file** at a fixed, project-declared standard path (declared once per §11.4.35, never moved without a §11.4.66 operator decision) — the OUT-OF-THE-BOX entry point for any fresh session (creating a new session = ONLY point the new agent at this one file). §11.4.131 promotes §11.4.127 (PREPARE-on-demand) into a STANDING, version-controlled ARTIFACT, ALWAYS present + ALWAYS in sync. (A) Exists at the declared path at all times (literal path per §11.4.35, never silently moved). (B) (Re)written whenever a fresh session is needed OR the live state materially changes (new HEAD / build-artifact id / phase / device state / in-flight job / blocking decision) — the §12.10 trigger set; a stale file = §11.4.131 violation (same class as a §12.10 stale-CONTINUATION). (C) Carries SHORT + FULL §11.4.127 variants; points to .remember/remember.md + docs/CONTINUATION.md read FIRST + git fetch ; embeds exact live-state anchors (HEAD, artifact checksums, device serials, in-flight PIDs+logs, evidence paths); states PHASE + NEXT + terminal goal; restates binding constraints (anti-bluff §11.4, no-force-push §11.4.113, exact version/naming, hardware gotchas); MOMENT-VALID never generic (§11.4.6). (D) §11.4.65 export (.html / .pdf siblings) + §11.4.44 revision header. (E) Given ONLY this file’s path a fresh session fully resumes with zero additional context. Composes

§12.10/.127/.65/.44/.6/.66/.126. Classification: universal (§11.4.17). Gate CM-COVENANT-114-131-PROPAGATION (literal 11.4.131) + rec. gate CM-SESSION-RESUMPTION-FILE-PRESENT + paired §1.1 mutation.

Canonical authority: Constitution.md §11.4.131. Release blocker. No

11 4 132
--skip-resumption-file / --ephemeral-prompt-only / --stale-resumption-OK / --generic-template-OK escape.
2026 06 07

§ . . — Risk-ordered validation
priority mandate (User mandate,
— —)

Forensic anchor — verbatim user mandate (2026-06-07): “We MUST ALWAYS first test and validate features, functionalities and fixes/changes that have been worked most recently, the ones which were most problematic, which have the most chance to crash or break again, the ones which have been re-opened the most times! Then, after we validate and verify all this with real (physical) proofs and hard evidence, with no false results and bluffs of any kind, we continue with all other existing tests in the test suites! This IS MANDATORY.”

Tests / validations / verifications MUST run in RISK-DESCENDING order — the highest-risk set FIRST, and ONLY AFTER that set is fully GREEN with real (physical) captured evidence does the remainder of the suite run. Closed risk factors, highest-first: (a) most-recently-worked; (b) historically most-problematic (longest defect history, most prior fixes/failures); (c) highest crash/break/regress likelihood (greatest blast radius/complexity/dependency surface); (d) most-reopened per §11.4.55 reopens-count. Each highest-risk item MUST pass with real (physical) captured evidence per §11.4.5/§11.4.69/§11.4.107 — no metadata-only / config-only / absence-of-error / grep-without-runtime PASS (§11.4/§11.4.1), no false results, no bluff (§11.4.6). ONLY AFTER the entire highest-risk set is GREEN with captured proof does the rest of the suite run; arbitrary-order or lower-risk-before-highest-risk-GREEN = §11.4.132 violation. REFINES/STRENGTHENS §11.4.130 (generalises “validate-just-fixed-first” to the full risk-ordered set) + §11.4.46 (adds explicit risk-ordering) + §11.4.42 (applies implementation-priority discipline to VALIDATION ordering). Classification: universal (§11.4.17) — consuming project supplies its recency/problematic-history/reopen-count sources (e.g. §11.4.93 DB reopens_count + last_modified) per §11.4.35. Composes

§11.4.4/.5/.6/.7/.40/.42/.46/.50/.55/.69/.107/.130. Gate CM-COVENANT-114-132-PROPAGATION (literal 11.4.132) + rec. gate CM-RISK-ORDERED-VALIDATION-PRIORITY + paired §1.1 mutation.

^{1 1 4 1 3 3}
Canonical authority: Constitution.md §11.4.132. Release blocker. No
--skip-risk-ordering / --any-order-OK / --suite-order-fixed escape.
_{2 0 2 6 0 6 0 8}

§ . . — Target-System + hardware
safety mandate (User mandate,
— —)

Forensic anchor — verbatim user mandate (2026-06-08): “Make sure that all changes we do to the System are ALWAYS safe for the System itself and for the hardware the system runs on! This is MANDATORY.”

Every change to the TARGET system (firmware, kernel, init/boot scripts, drivers, sysfs/devfreq/voltage/clock/thermal/regulator register writes, partition/bootloader/U-Boot, HAL, framework, prebuilts, device config) MUST ALWAYS be safe for BOTH (a) the target System itself — MUST NOT brick, boot-loop, corrupt data, or render the device unrecoverable — AND (b) the hardware it runs on — MUST NOT exceed safe electrical/thermal/voltage/clock limits or damage panels/storage/radios/regulators. Concrete obligations: (1) reversible-first — verify irreversible high-blast-radius changes (bootloader/U-Boot MD5, partition layout) against known-good values + capture a pre-op backup (§9.2) BEFORE applying; (2) NO unverified hardware-control writes — never write an unverified value to a voltage/clock/regulator/thermal-throttle/current-limit sysfs node or register that could exceed datasheet limits, the safe range established as FACT (§11.4.6), never guessed; (3) thermal/perf changes (forcing a performance governor, pinning the top OPP, disabling thermal management) MUST respect the device’s cooling design, validated by captured thermal evidence; (4) flashing MUST use the sanctioned tool + a freshly-built integrity-verified image — never an ad-hoc partition write or stale/unverified artifact; (5) unprovable-safety ⇒ blocked — a change whose target/hardware safety cannot be established from captured evidence is treated as UNSAFE and blocked (§11.4.6 + §11.4.101 reversible-first + §11.4.123 rock-solid-proof). DISTINCT from §12 host-session safety: §12 protects the DEVELOPER’s HOST + session; §11.4.133 protects the TARGET device + its hardware — both apply, neither weakens the other. Classification: universal (§11.4.17) — consuming project supplies its concrete hardware-control surfaces, datasheet-safe ranges,

known-good bootloader/image hashes, and sanctioned flashing tool per §11.4.35. Composes §12 / §11.4.6 / §11.4.101 / §11.4.108 / §11.4.123. Gate CM-COVENANT-114-133-PROPAGATION (literal 11.4.133) + rec. gate CM-TARGET-HARDWARE-SAFETY + paired §1.1 mutation.

Canonical authority: Constitution.md §11.4.133. Release blocker. No `--unsafe-hardware-write` / `--skip-system-safety` / `--brick-risk-accepted` escape.

§ . . — Code-review iterate-until-GO + rock-solid-evidence mandate (User mandate, — —)

Forensic anchor — verbatim user mandate (2026-06-08): “For any fixes/changes given back to us for re-work by the code-review process, once we fix/improve everything per the code-review’s requests, we MUST RE-RUN code-review AGAIN until we get a GO from it with NO new issues reported or warnings of any kind! All results produced by this whole process MUST ALWAYS give us rock-solid PHYSICAL evidence that the fixed/improved codebase really works now as expected, with no false results and no bluff(s) of any kind.”

When the §11.4.125 code-review returns ANY finding (BLOCKING, nit, or warning) and the author fixes/improves per that review, the code review MUST BE RE-RUN — and KEEP being re-run after each remediation round — until a clean GO with ZERO new issues AND ZERO warnings. A single “addressed the findings” pass is NOT sufficient: the corrected batch MUST pass a FRESH adversarial review (a re-review can surface NEW findings introduced by the fixes — the §11.4.1 fix-A-creates-B mode). The loop terminates ONLY on a clean GO; a residual warning is itself a finding that re-arms the loop. Every round’s verdict AND every fix’s validation MUST carry rock-solid PHYSICAL captured evidence per §11.4.5 / §11.4.69 / §11.4.107 (captured audio / video / sysfs / dumphsys / sink-side / runtime-signature) proving the fixed codebase REALLY works — never metadata-only / config-only / absence-of-error / grep-without-runtime; no false results, no bluff; a GO unbacked by physical evidence is a §11.4 PASS-bluff at the review-loop layer. REFINES / STRENGTHENS §11.4.125 (iterate-until-no-blocking): makes the loop EXPLICIT, raises termination to ZERO findings AND ZERO warnings, BINDS physical evidence to every round. Classification: universal (§11.4.17). Composes §11.4.125 /

§11.4.1 / §11.4.4 / §11.4.5 / §11.4.6 / §11.4.69 / §11.4.107 / §11.4.50 / §11.4.108 / §11.4.123. Gate `CM-COVENANT-114-134-PROPAGATION` (literal `11.4.134`) + rec. gate `CM-CODE-REVIEW-ITERATE-UNTIL-GO` + paired §1.1 mutation.

Canonical authority: `Constitution.md` §11.4.134. Release blocker. No `--skip-rereview` / `--single-review-pass` / `--warnings-ok` / `--evidence-optional` escape.

§11.4.135 — Standing regression-guard suite + every-fixed-defect-gets-a-permanent-regression-test (User mandate, 2026-06-08). Every project MUST maintain a STANDING regression-guard suite that runs on EVERY build+deploy and BLOCKS the release tag on any failure. Every closed defect (stable ticket id, e.g. ATM-NNN) MUST, in the SAME commit as its fix (extending the §11.4.43 DOCUMENT step), register a permanent §11.4.115 RED-on-broken-artifact regression test into the suite — `RED_MODE=1` capturing the historical defect on a pre-fix artifact (the proof the guard is real), `RED_MODE=0` the standing GREEN guard asserting the defect is ABSENT. A closure without a registered guard is a §11.4.123 violation. The suite runs FIRST in the post-deploy cycle (highest-risk set per §11.4.132) and is a §11.4.40 release-gate blocker. Forensic anchor (FACT): the wrong-subtitle-on-2nd-display defect was “fixed” via a source-side `CONTROL_MENU_LABEL_DENYLIST` that NO test mirrored or re-ran, so the NEXT chrome class recurred silently while the GREEN suite passed. Industry-standard bug-driven testing (Google content-driven testing; AOSP CTS/Tradefed) made mechanical + enforced. Composes §11.4.4 / §11.4.40 / §11.4.43 / §11.4.46 / §11.4.50 / §11.4.107 / §11.4.108 / §11.4.115 / §11.4.118 / §11.4.123 / §11.4.124 / §11.4.130 / §11.4.132. Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-135-PROPAGATION` (literal `11.4.135`) + recommended gates `CM-REGRESSION-GUARD-REGISTERED` / `CM-REGRESSION-GUARD-SUITE-WIRED` + paired §1.1 mutation. **Canonical authority:** `constitution submodule Constitution.md` §11.4.135. Non-compliance is a release blocker. No escape hatch — no `--skip-regression-guard`, `--no-guard-on-close`, `--guard-optional` flag.

§11.4.136 — Real-content end-to-end playback-test mandate (User mandate, 2026-06-08). Refines/strengthens §11.4.107. Any test asserting media playback works MUST drive REAL content (catalog stream or offline reference clip) through the user’s path (§11.4.48 UI-driven → §11.4.117 CV/OCR fallback) and assert it genuinely PLAYS via the §11.4.107 liveness battery PLUS a decoder-health census — a numeric drop-buffer budget, no buffer-timestamp re-order/discard, no codec-

reject (cite Android/Media3 ExoPlayer OEM pre-OTA playback-test mandate: “too many dropped buffers” >25, “unexpected presentation timestamp”, “test timed out”). Metadata-only / launch-only / registration-only / single-frame / config-only PASS is forbidden (§11.4 / §11.4.1). A golden/reference clip corpus (BBC ExoPlayer testing samples) is the offline ground-truth. Composes §11.4.5 / §11.4.48 / §11.4.50 / §11.4.107 / §11.4.117 / §11.4.123 / §11.4.13 / §11.4.69. Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-136-PROPAGATION` (literal `11.4.136`) + recommended gate `CM-REAL-CONTENT-PLAYBACK-TEST` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.136. Non-compliance is a release blocker. No escape hatch — no `--launch-proves-playback`, `--skip-decoder-health`, `--metadata-playback-pass-suffices` flag.

§11.4.137 — Subtitle/caption content-correctness oracle + secure-display-proxy-honesty mandate (User mandate, 2026-06-08). Refines §11.4.117 + §11.4.107 + §11.4.112. Forensic anchor (FACT): tests tasked to “physically verify the 2nd-display subtitle” PASSEd GREEN while subtitles did NOT show / showed WRONG — the streaming player surface is `FLAG_SECURE` so `screencap -d <secondary>` returns BLACK (autonomous PIXEL verification structurally impossible per §11.4.112), so the test fell back to the accessibility-scraped/`persist.atmosphere.subdebug` proxy, and the proxy accepted a chrome/menu LABEL (`Аудио и субтитры`) as a valid subtitle because the prose floor accepted any multibyte prose and NO menu-label denylist + NO position/cadence check existed. The mandate: a subtitle-correctness test MUST classify the cue’s *content class* — a present cue is NOT a correct cue. CHROME (FAIL) if a known control/menu label (closed multilingual deny-list MIRRORED from source, case-folded incl. non-ASCII), time/numeric chrome, not prose, OUTSIDE the lower safe-title band (CEA-708 9-anchor grid), OR STATIC across the window (real subtitle changes → ≥2 distinct prose cues, a metamorphic relation). DIALOGUE (PASS) only when prose + not-denied + not-chrome + position-ok + cadence ≥2 OR fuzzy-matches the SOURCE-extracted cue via normalized edit distance (§11.4.123 host ground truth). The oracle MUST be self-validated golden-good/golden-bad (§11.4.107(10)) and the deny-list MUST be verified present in the SHIPPED artifact (§11.4.108) — a source-green denylist with no test mirror + no artifact check is the exact recurrence pattern forbidden here. Secure-display honesty (§11.4.112): where `FLAG_SECURE` makes pixel verification impossible, the rock-solid autonomous proof is the player’s caption telemetry + source-track presence + content-class

oracle — NEVER a faked pixel “physical” pass; human-eye pixel confirmation is `operator_attended` (§11.4.52) with a tracked migration item. App-agnostic (keys off content class). Composes §11.4.3 / §11.4.5 / §11.4.6 / §11.4.107 / §11.4.108 / §11.4.112 / §11.4.115 / §11.4.117 / §11.4.123 / §11.4.13 / §11.4.69. Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-137-PROPAGATION` (literal `11.4.137`) + recommended gate `CM-SUBTITLE-CONTENT-CORRECTNESS-ORACLE` + paired §1.1 mutation (strip the denylist/position/cadence check → golden-bad `Аудио и субтитры` `PASSes` → gate `FAILs`). **Canonical authority:** constitution submodule `Constitution.md` §11.4.137. Non-compliance is a release blocker. No escape hatch — no `--present-cue-is-correct`, `--skip-chrome-oracle`, `--length-heuristic-suffices`, `--pixel-pass-on-secure-display`, `--skip-position-check`, `--skip-cadence-check` flag.

§11.4.138 — Operator-escape => mandatory bluff-audit + permanent guard (User mandate, 2026-06-08). When the operator (or any out-of-band channel) finds a defect that the GREEN test suite passed, this is by definition a §11.4 `PASS-bluff` — it MUST trigger, before the fix is closed: (1) a §11.4.102 systematic-debugging pass to `FACT-root-cause`; (2) a bluff-audit identifying the EXACT assertion that should have caught it but didn’t, cited to `file:line` (canonical example: `lib/subtitle_content_validation.sh:sub_is_prose()` returning `TRUE` for `Аудио и субтитры`); (3) a permanent §11.4.135 regression guard registered in the SAME commit as the fix, with its §11.4.115 `RED` capturing the operator-found defect; (4) the bluff-audit committed under `docs/research/<scope>/<defect>_bluff_audit/`. Closing an operator-found defect WITHOUT the bluff-audit + permanent guard is itself a §11.4 violation (the bluff that let it through is still live and the defect will recur). Composes §11.4 / §11.4.1 / §11.4.102 / §11.4.108 / §11.4.115 / §11.4.118 / §11.4.123 / §11.4.135. Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-138-PROPAGATION` (literal `11.4.138`) + recommended gate `CM-OPERATOR-ESCAPE-BLUFF-AUDIT` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.138. Non-compliance is a release blocker. No escape hatch — no `--close-without-bluff-audit`, `--operator-find-is-just-a-bug`, `--skip-permanent-guard` flag.

§11.4.139 — Fresh-process clean-artifact runtime-signature mandate (User mandate, 2026-06-08). Refines §11.4.108. Before any post-deploy validation — ESPECIALLY a non-pixel proxy verification (the subdebug/accessibility-cue channel used for `FLAG_SECURE` displays) — the harness MUST assert `running-artifact ==`

built-artifact: the deploy yielded a CLEAN target (mutable-overlay/userdata wiped) OR a pre-validation check proves no stale overlay shadows the deployed code (e.g. every guarded package — incl. the Presenter that emits the subtitle cue — resolves to the system partition, no per-user override). A stale shadow of the cue-emitting component (e.g. a Presenter APK predating the denylist) makes the proxy report on code that was never deployed — any PASS is a §11.4 PASS-bluff. Each fix declares ONE machine-checkable runtime signature verified on the clean target (the §11.4.108 registry IS the definition of done); for the subtitle class the signature is “the shipped Presenter APK contains the denylist literal (case-insensitive) AND the subdebug channel emits `candidate REJECTED reason=chrome-label` for a menu label.” Composes §11.4.46 / §11.4.108 / §11.4.130 / §11.4.135 / §11.4.137. Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-139-PROPAGATION` (literal `11.4.139`) + recommended gate `CM-CLEAN-ARTIFACT-RUNTIME-SIGNATURE` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.139. Non-compliance is a release blocker. No escape hatch — no `--validate-against-running-state`, `--skip-clean-precondition`, `--shadow-OK` flag.

§11.4.140 — Universal action-prefix system (ACTION_NAME ::) (User mandate, 2026-06-09; GRAMMAR_ADDENDUM 2026-06-09; arrow form 2026-07-02). When a user prompt's FIRST non-blank line starts with a recognised action prefix, you MUST: (1) look the action token up in the action registry `constitution/actions/registry.yaml` (or `$HELIX_ACTION_REGISTRY`); (2) if it is a registered action, REPLACE the prefix with that action's `expansion` text and apply its `rules` ; (3) execute the remainder of the prompt under the expanded instruction. **Five EQUIVALENT forms** — same action, same expansion, same execution: (1) `ACTION_NAME :: <rest>` (bare `::`), (2) `PREFIX::ACTION_NAME :: <rest>` (namespaced `::`), (3) `/ACTION_NAME <rest>` (bare slash), (4) `/PREFIX::ACTION_NAME <rest>` (namespaced slash), (5) `ACTION_NAME ---> <rest>` (bare arrow) $\equiv$ `PREFIX::ACTION_NAME ---> <rest>` (namespaced arrow). Thus `BACKGROUND :: x` $\equiv$ `DEFAULT::BACKGROUND :: x` $\equiv$ `/BACKGROUND x` $\equiv$ `/DEFAULT::BACKGROUND x` $\equiv$ `BACKGROUND ---> x` $\equiv$ `DEFAULT::BACKGROUND ---> x` . `PREFIX` is an action NAMESPACE; the reserved default namespace is **DEFAULT** , and an action runs WITH or WITHOUT the prefix. Grammar (all hold): anchored at the FIRST non-blank line only (mid-prose tokens never match); the action token AND the namespace are UPPERCASE-only `[A-Z][A-Z0-9_]*` (lowercase never matches); the namespace separator `::` inside the token carries NO surrounding spaces (`PREFIX::ACTION_NAME`), DISTINCT from the action-body separator `" :: "` (one ASCII space on each side of `::` — avoids `C+ + Foo::Bar` , YAML `key: value` , URLs) in forms 1/2, the arrow-body separator `" ---> "` (one ASCII space on each side) in form 5, and the slash-body separator (one space) in forms 3/4; stacked prefixes (`A :: B :: rest`) apply outer-to-inner, left-to-right (expand `A` , re-scan, expand `B` , then the residual is the task); a leading `\` escapes the prefix for the `::` , arrow, and slash forms (`\BACKGROUND :: x` , `\BACKGROUND ---> x` , `\ /BACKGROUND x` — treat literally, strip the backslash, NO expansion) so action names can be discussed. **Conflict rule (slash form):** `/ACTION_NAME` (form 3) is honored as the action ONLY when `ACTION_NAME` (case-folded) does not collide with a built-in/host slash command (registry `slash_bare: auto` + `slash_conflicts: [...]`); form 4 (`/PREFIX::ACTION_NAME`) is ALWAYS unambiguous and always honored. An unknown token that matches the grammar shape (any of the 5 forms) but is NOT registered is NEVER silently

expanded or silently dropped — ask which registered action was meant (§11.4.66 / §11.4.105) or treat it as a literal prompt, NEVER invent an expansion (§11.4.6); any prompt not satisfying the grammar is an ordinary prompt and the system is a no-op. The registered action **BACKGROUND** expands to: *“The following prompt that we will provide MUST BE executed in background in parallel with all main work streams using the subagents-driven development approach! All work done MUST PRODUCE rock solid evidence covered with hard physical proof(s) that all done is working as expected and as specified without any false results and without any bluff!”* (composes §11.4.20 / §11.4.70 subagent-driven, §11.4.58 / §11.4.103 parallel streams, §11.4.89 background execution, §11.4.5 / §11.4.69 / §11.4.107 captured physical evidence, §11.4 anti-bluff). A second registered action **REMINDER** re-surfaces previously-scheduled, CRITICAL, status-UNCERTAIN work: FIRST verify the ACTUAL current status from captured evidence (never assume done or not-done — §11.4.6 no-guessing), THEN act on the delta — report the captured proof if genuinely complete, resume from the exact point if partial, action it NOW if not started, or surface the block per §11.4.66/ §11.4.101 — always producing a status verdict (composes §11.4.6 / §11.4.87 / §11.4.94 / §11.4.97 / §11.4.103 / §11.4.108 / §11.4.130 / §11.4.147).

Severity / handling-priority markers (§11.4.140 extension, 2026-07-16).

Three further registered actions tag the REMAINDER of the prompt with a HANDLING PRIORITY (as **BACKGROUND** tags it “run in background”):

CRITICAL (highest-priority / potentially release-blocking — maximum urgency + rigor ahead of lower-priority work, track it, auto-activate §11.4.102 systematic-debugging when an issue is involved, do NOT defer it),

IMPORTANT (high-priority, above routine work but below CRITICAL — track it, full anti-bluff rigor), and **NOTE** (capture the remainder as durable

CONTEXT — request-history §11.4.208 / §11.4.210 + persistent memory when a non-obvious project fact — applied when relevant, NOT an urgent action unless it contains an explicit action; §11.4.6 record only what the note says, never invent). All six grammar forms work (**CRITICAL**: x ≡

CRITICAL :: x ≡ **CRITICAL** ---> x ≡ /**CRITICAL** x ≡ /

DEFAULT::**CRITICAL** x). Registering **NOTE** promotes **NOTE**: from an ordinary-prose NO-OP to a registered action (ordinary-prose single-colon examples are now **TODO**: / **WARNING**: / **FIXME**:).

The system is UNIVERSAL (every CLI agent reads this block via its context carrier per §11.4.35), extensible (new action = new registry row), decoupled + reusable (§11.4.28), and loads out-of-the-box. Classification: universal (§11.4.17). **Canonical authority:** constitution submodule `Constitution.md` §11.4.140. Non-compliance is a release blocker. No escape hatch — no `--skip-action-prefix`, `--ignore-prefix`, `--no-registry`, `--invent-expansion-OK`, `--single-layer-only` flag.

§11.4.141 — Token-efficiency mandate (research-derived + operator mandate, 2026-06-09). Every project worked on by AI coding agents **MUST** cut token spend (input AND output) toward **30–40% of current (a 60–70% reduction)** WITHOUT degrading quality/performance/safety or breaking any existing mechanism, via a composable, safety-ranked measure set: (1) **prompt-cache the static governance prefix** — the always-loaded governance forms a byte-stable cache-breakpointed prefix with no volatile bytes ahead of it; cache reads cost $\sim 0.1 \times$ base input (the dominant cost driver — measured $\sim 170\text{K}$ tokens of governance re-sent every turn, externally corroborated by Claude Code issue #24147); caching is transparent so it removes no rule, weakens no gate, changes no verdict — only billing (PRIMARY, biggest + safest lever); (2) **subagent model-tiering + output-to-file** — mechanical non-judgment work (search/grep/status/doc-export/read-only probes) to a Haiku-class model, the strong model RESERVED for all reasoning/verdicts/fix-design (§11.4.102)/code-review (§11.4.125)/demotion (§11.4.7), large output persisted to a file not an inline 350–520K-token transcript; the cheap model never emits a PASS so §11.4.50 + anti-bluff are untouched; (3) **thin always-loaded INDEX + on-demand detail** — concise index (one line per fix/anchor, EACH carrying the literal `11.4.N` token so propagation gates pass) with the canonical full text kept gate-scanned in `constitution/Constitution.md` and reachable in one hop — a deduplication realising §11.4.35, never a deletion; (4) **CodeGraph/retrieval-first over full-file loading** (§11.4.78/§11.4.79); (5) **output-token reduction** — terse status + `effort:"low"` on the mechanical allowlist only; (6) **tool-call batching + no re-reads**; (7) **compaction/context-editing for long sessions. Mandatory measured proof:** a token-accounting harness measures tokens-per-development-cycle BEFORE vs AFTER on a frozen deterministic workload from the authoritative `usage` object (input/cache_read/cache_creation/output split; NEVER `tiktoken`, NEVER the client-side cost estimate), reproduced N times (§11.4.50), pass = AFTER $\leq 40\%$ of BEFORE OR the measured best-safe reduction with a cited cold-

cache reason; the AFTER run MUST show ZERO regression on the pre-build sweep + meta-test mutation sweep + propagation gates + a strong-model reasoning probe + a cache-warm proof (`cache_read_input_tokens > 0`) — cost reduction with quality regression is a §11.4 FAIL. The headline number is the *measured* reduction, never the design estimate (§11.4.6/§11.4.123). No measure may break/degrade any existing mechanism, and the rule is structured so none can. Composes

§11.4.5/.6/.20/.40/.50/.58/.69/.70/.78/.79/.80/.103/.106/.123/.125/.128/§12.6/§1.1.

Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-141-PROPAGATION` (literal `11.4.141`) + recommended gate `CM-TOKEN-EFFICIENCY` + paired §1.1 mutation (inject a pre-breakpoint volatile token → cache collapses → measured reduction falls below bar → gate FAILs). **Canonical authority:**

constitution submodule `Constitution.md` §11.4.141. **Project instantiation**

(§11.4.35): [the consuming project fills in: confirm the Claude Code governance-prefix cache is warm + free of pre-breakpoint volatiles; tier the §11.4.96-SAFE mechanical subagents (search/grep/status/doc-export) to Haiku with output-to-file under `qa-results/` ; replace the consumer `CLAUDE.md` 112-row Applied-Fixes inline table + verbatim anchor restatements with a literal-anchor index pointing at `constitution/Constitution.md` ; keep CodeGraph/lumen-first navigation; run the harness BEFORE/AFTER to prove the measured reduction.] Non-compliance is a release blocker. No escape hatch — no `--skip-token-efficiency` , `--no-cache-governance` , `--assert-reduction-without-measuring` , `--tier-down-reasoning` , `--inline-all-governance` , `--tiktoken-estimate-OK` flag.

§11.4.142 — Universal code-review mandate — every change reviewed, always, no exception (User mandate, 2026-06-09). Verbatim operator mandate:

“ALL changes we do MUST pass through the code review step!!! ALWAYS!!!”

EVERY change made to ANY repository this Constitution governs — without exception — MUST pass through an INDEPENDENT code-review step BEFORE it is accepted, committed, or built. NO change class is exempt: source code, fixes, tests, gates, meta-test mutations, documentation, doc-tooling, build/CI scripts, configuration, governance files (Constitution / CLAUDE / AGENTS / QWEN), conductor main-stream edits, sub-agent output, refactors, single-line edits — if a diff exists, it gets an independent review. This is the ABSOLUTE form of §11.4.125 (code-review agent gate after a batch, before pre-build sweep + main build):

§11.4.142 strips every implicit scoping §11.4.125’s “after all fixes/changes/ implementations are done” phrasing could leave open — no “just a doc edit”, no

“just a one-liner”, no “the author already self-reviewed (§11.4.92)”, no “trivial change” carve-out. **Independence is load-bearing** — the reviewer MUST be structurally separated from the author (a dedicated code-review agent, subagent-driven by default per §11.4.70 / §11.4.20, or a distinct human), NEVER the author re-reading their own work; §11.4.92’s multi-pass self-evaluation is the AUTHOR-side discipline and PRECEDES (never satisfies) §11.4.142. **The review is itself anti-bluff** (§11.4 / §11.4.1) — a rubber-stamp “looks good” is not a review; it MUST genuinely analyse correctness, safety (no host §12 / data §9 / target-hardware §11.4.133 regression), will-it-really-work (no solve-A-create-B), end-user behaviour (§11.4 / §107), and test-genuineness (§1.1), its conclusions captured evidence per §11.4.5 / §11.4.69, and **it iterates to a clean GO per §11.4.134** — any finding (BLOCKING / nit / warning) re-arms the loop, acceptance only on ZERO new findings + ZERO warnings with rock-solid physical evidence. Honest boundary (§11.4.6): a passing review does NOT replace §11.4.108 four-layer runtime-signature verification nor the §11.4.40 full-suite retest — it is one of MULTIPLE STRONG LAYERS, and the FIRST one every change crosses. Composes §11.4.1 / §11.4.4 / §11.4.5 / §11.4.6 / §11.4.20 / §11.4.40 / §11.4.69 / §11.4.70 / §11.4.92 / §11.4.108 / §11.4.110 / §11.4.125 / §11.4.134 / §107 / §1.1. Classification: universal (§11.4.17) — the consuming project supplies its reviewer-dispatch mechanism + change-acceptance seam (commit wrapper / merge queue / PR gate) per §11.4.35. Propagation gate `CM-COVENANT-114-142-PROPAGATION` (literal `11.4.142`) + recommended gate `CM-EVERY-CHANGE-REVIEWED` (every accepted change carries a fresh independent-review-completed marker for its diff, produced by a reviewer distinct from the author, before acceptance/commit/build) + paired §1.1 mutation (strip the literal → propagation gate FAILs; accept a diff with no independent-review marker → `CM-EVERY-CHANGE-REVIEWED` FAILs). **Canonical authority:** constitution submodule `Constitution.md` §11.4.142. Non-compliance is a release blocker. No escape hatch — no `--skip-review`, `--no-review`, `--trivial-change-exempt`, `--doc-edit-exempt`, `--self-review-suffices`, `--review-after-commit` flag.

§11.4.143 — Real-user-journey mandate for video-streaming-app full-automation tests (User mandate, 2026-06-10). Verbatim operator mandate: “All video streaming apps ... require to choose some title and to press proper UI button to start or resume playing! Proper UI interaction to play exact show with proper content and subtitles is MANDATORY! Without it we just eventually play on 2nd display sample, and that’s it mostly! THIS MUST BE ADDED as MANDATORY

RULE regarding testing of any video streaming app in general with full automation tests! Universal in root constitution + a consuming project's extensions." Any full-automation test that asserts a video player / streaming application plays content MUST drive the REAL end-user journey through the app's OWN UI — launch → BROWSE the actual catalog → choose a SPECIFIC title → press the real Play/Resume button → confirm THAT chosen content is genuinely playing, with its correct subtitles, on the intended routing target. A test that bypasses the journey with a sample / demo / built-in-loop clip, a deep-link / `am start -a VIEW` / intent shortcut, a synthetic or pre-staged stream, or any path that does NOT exercise the app's own browse-select-play UI is a §11.4 PASS-bluff at the user-journey layer: it validates ROUTING (that *something* reaches the display) while leaving the user-visible behaviour the operator cares about — “the show I picked actually plays” — unproven (the operator's “we just eventually play on 2nd display sample, and that's it mostly” gap). The mandate (ALL hold): (1) **Real journey, not a shortcut** — launch → catalog browse → specific-title selection → real Play/Resume press through the app's own UI (§11.4.48 UI-driven), NEVER a deep-link / intent / sample / loop-clip shortcut; (2) **Chosen-content confirmation** — the PASS proves the SPECIFIC selected title is playing (not merely that pixels move on the target) via the §11.4.107 liveness battery on the §11.4.136 real-content path + the §11.4.137 subtitle content-correctness oracle for the chosen title's captions; (3) **Non-introspectable UIs use the pixel oracle** — when the app's accessibility hierarchy is blank/unreliable (TV-Compose / leanback / canvas / GL), DRIVE input + ASSERT content via the §11.4.117 CV/OCR pixel oracle, never a hierarchy-only tool; (4) **Login via the credential single-source** (§11.4.10), never hardcoded, never logged; (5) **Honest SKIP, never a faked PASS** — where the autonomous journey is genuinely infeasible (hard human-only login / CAPTCHA, geo-block per §11.4.3, secure-surface pixel-blanking per §11.4.112) the test is `operator_attended SKIP-with-reason` per §11.4.52 + §11.4.3 with a tracked migration item — NEVER a metadata-only / sample-played / routing-only PASS. Honest boundary (§11.4.6): “the routing fired so the title is playing” is a guess — only the chosen-content liveness + subtitle oracle on the real journey proves it. Composes §11.4.48 / §11.4.107 / §11.4.117 / §11.4.136 / §11.4.137 / §11.4.52 / §11.4.3 / §107 / §1.1. Classification: universal (§11.4.17) — the consuming project supplies its concrete app roster, login-credential source, routing target, and UI-driving / pixel-oracle harness per §11.4.35. Propagation gate `CM-COVENANT-114-143-PROPAGATION` (literal `11.4.143`) + recommended gate `CM-VIDEO-REAL-JOURNEY-TEST` (every video-streaming-app playback test drives the

real browse-select-play journey + confirms the chosen content + subtitles, or SKIPS-with-reason) + paired §1.1 mutation (replace a real-journey test's browse-select-play path with a deep-link / sample shortcut → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.143. Non-compliance is a release blocker. No escape hatch — no `--sample-playback-ok` , `--skip-real-journey` , `--deep-link-suffices` , `--routing-only-pass` , `--no-title-selection` flag.

§11.4.144 — Tracked/recorded-device availability-following mandate (User mandate, 2026-06-10). Direct operator mandate: every device the project is tracking / following / recording (test / debug / manual-testing device, across every reachable transport — USB / wireless ADB / SSH / serial / network introspection API) MUST be availability-FOLLOWED — its connection state continuously monitored, any drop handled, never silently abandoned and never presented as a continuous recording. The §11.4.128 always-on recorder KNOWS a tracked device is absent (its per-device loop guards on the reachable state) but, lacking a following discipline, merely spins idle — no data captured, no offline event logged, no resume, no escalation — so the recording corpus presents a continuous timeline with a silent hole, a §11.4 PASS-bluff at the recording-integrity layer (it claims continuous capture while no data exists for the gap). On a tracked device leaving its reachable state the system MUST, automatically: (a) **DETECT** the drop and **log an honest offline event** into the recording corpus (§11.4.6 — a silent gap presented as continuous capture is a fabricated-continuity bluff; §11.4.128 — a silent recording gap defeats always-on recording); (b) **WAIT** for the device to return using the project's ALREADY-DEFINED reconnection timings — never invented numbers (§11.4.6), the SAME grace / reconnect / poll budgets the project's recovery path already uses; (c) **RE-ATTACH** — resume recording / tracking the moment the device returns and log an honest online / resume event; (d) **ESCALATE** to the project's sanctioned device-recovery path (per §11.4.69 feature class `device_recovery`) if the device does not return within the defined timeout — through the sanctioned, authorization-gated recovery entry point ONLY, never bypassing its gate and never performing a destructive recovery (e.g. a power-cycle) autonomously without that authorization (§11.4.21 / §11.4.101 — high-blast-radius recovery is gated; while blocked the system keeps following the device, logging the blocked-escalation honestly). A tracked-device drop that produces a silent corpus hole, a never-resumed recording, or a never-escalated permanent absence is the

bluff this anchor forbids. Composes §11.4.128 (always-on recording — §11.4.144 closes its drop-handling gap) / §11.4.69 (`device_recovery` sink-side positive evidence) / §11.4.6 (honest offline / online events, reused-not-invented timings) / §11.4.14 (watchdog children reaped on stop) / §11.4.21 + §11.4.101 (gated, non-autonomous escalation). Classification: universal (§11.4.17) — the consuming project supplies its concrete tracking transport, device set, already-defined reconnection timings, and sanctioned recovery entry point per §11.4.35.

Propagation gate `CM-COVENANT-114-144-PROPAGATION` (literal `11.4.144`) + recommended gate `CM-DEVICE-AVAILABILITY-FOLLOWED` + paired §1.1 meta-test mutation (strip the watchdog wiring / let an absence go unlogged → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.144. Non-compliance is a release blocker. No escape hatch — no `--skip-availability-following` , `--silent-recording-gap-OK` , `--no-reconnect-wait` , `--invent-reconnect-timing` , `--skip-recovery-escalation` , `--autonomous-power-cycle-OK` flag.

§11.4.145 — Independent multi-angle impact-research per change (User mandate, 2026-06-10). For EVERY fix / change / new feature, INDEPENDENT impact-research agents (subagent-driven §11.4.70/§11.4.20, structurally separate from the author — §11.4.92 self-eval PRECEDES, never satisfies — adversarial “refute-the-change” stance to defeat the documented LLM confirmation-bias failure mode) MUST research the change AND its connected/dependent features across a CLOSED SET OF EIGHT ANGLES — (1) correctness/logic, (2) regression (what existing working feature could break — via call-graph impact enumerating every direct + transitive caller and contract-dependent feature, each mapped to its test), (3) latent/dangerous-code — the recents class (runtime-only failure, interface/ABI/contract mismatch, concurrency/data-race, resource lifecycle), (4) security (taint/injection/secret/unsafe-API), (5) performance (latency/memory/thermal under load vs baseline), (6) host/data/target-hardware safety per §12/§9/§11.4.133, (7) cross-feature interaction (shared state/timing/hardware contention), (8) business-logic conformance (matches the spec AND the connected features’ contracts, not merely compiles). Forensic anchor (FACT, project-internal): the recents / 3-button-nav path shipped a latent AIDL-interface-contract mismatch that PASSED code review yet failed at RUNTIME ONLY (ANR on a recents-reaping gesture) because no review angle interrogated the interface contract / concurrency-lifecycle / silently-broken connected feature — the §11.4.108 SOURCE → ARTIFACT → RUNTIME → USER-

VISIBLE gap caught only after the fact; §11.4.145 shifts that discovery LEFT. Each angle MUST name the TOOL(S) used + obtain the required DEPTH (the diff, the full changed unit, its declared contract, the spec section, every connected feature — NEVER the diff lines alone) + emit a captured-evidence conclusion per §11.4.5/§11.4.69 (“looks fine” forbidden, §11.4.1); a genuinely-N/A angle is recorded NOT_APPLICABLE: <reason> per §11.4.6, never silently skipped. Output = a per-change impact-research REPORT (one section per angle: tool + evidence path + verdict) + a single GO/NO-GO that BLOCKS acceptance/commit/build on ANY unmitigated risk in ANY angle; the change is fixed/mitigated + affected angles re-researched, iterating to a CLEAN GO (zero unmitigated risk + zero warning) per §11.4.134 before the §11.4.125 review gate. Honest boundary (§11.4.6): a GO proves cross-angle internal-consistency + bounded blast-radius — it does NOT replace §11.4.108 runtime-signature verification nor §11.4.40 full-suite retest; it is one of MULTIPLE STRONG LAYERS, the FIRST research pass every change crosses. Composes/strengthens §11.4.92/.125/.108/.110/.134/.78/.79/.107/.85/.133/§12/§9/.99/.6/§1.1. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-145-PROPAGATION (literal 11.4.145) + recommended gate CM-IMPACT-RESEARCH-PER-CHANGE + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; accept a change with a report missing an angle or an unmitigated NO-GO → CM-IMPACT-RESEARCH-PER-CHANGE FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule Constitution.md §11.4.145. Non-compliance is a release blocker. No escape hatch — no --skip-impact-research, --single-angle-suffices, --author-researches-own-change, --diff-skim-OK, --no-go-overrideable, --trivial-change-no-research flag.

§11.4.146 — Reproduce-first test + same-test-confirms-fix + mandatory extend-to-all-cases workflow (User mandate, 2026-06-10). Every reported problem MUST be handled by a NAMED three-step test workflow — §11.4.146 does NOT re-author its component disciplines, it NAMES + BINDS them into the operator’s exact per-defect sequence and adds ONLY the two emphases the components leave implicit. **(STEP 1 — REPRODUCE-FIRST + INVESTIGATE)** before any fix, author the §11.4.43 RED test as a §11.4.115 RED-baseline-on-the-broken-artifact (reproduce the defect on the CURRENT pre-fix artifact, capture defect-present physical evidence per §11.4.5/§11.4.69/§11.4.107); NEW EMPHASIS (D1) that same RED test is ALSO a deliberate investigation instrument per §11.4.102(A) — it MUST gather ADDITIONAL forensic data characterising the

defect (triggers, boundaries, input/topology scope, adjacent failure modes) that feeds BOTH the fix design AND the STEP-3 extend scope; a RED test used only as a binary present/absent gate with no characterisation satisfies §11.4.115 but NOT §11.4.146 STEP 1. **(STEP 2 — SAME-TEST-CONFIRMS-FIX)** after the fix the SAME test source (the §11.4.115 polarity switch flipped `RED_MODE=1→0`) confirms the defect ABSENT — RED-on-broken then GREEN-on-fixed, both captured, on a clean target per §11.4.108/§11.4.139, validated-first-after-redeploy per §11.4.130, deterministic per §11.4.50; no separate happy-path test substitutes (the §11.4.43/§11.4.115 PASS-bluff). **(STEP 3 — EXTEND-TO-ALL-CASES, mandatory per-fix)** NEW EMPHASIS (D2) immediately after STEP 2 the test MUST be FANNED OUT across the full case-space of the SAME functionality — all flows, valid + invalid + boundary (§11.4.85 empty/max/off-by-one) + concurrent (§11.4.85 contention) + failure-injection (§11.4.85 chaos) + topology variants (§11.4.3) — confirming no issue of any kind, with PROVABLE enumerated coverage per §11.4.118 (a listed case-set with per-case outcome, never “no other issues found”); a REQUIRED step, NOT deferred to a release-cycle discovery sweep and NOT reduced to a single guard; each user-visible case carries rock-solid physical evidence per §11.4.123 + is registered into the §11.4.135 standing regression-guard suite; a newly-discovered fan-out defect triggers §11.4.4 test-interrupt + re-enters STEP 1. (ANTI-BLUFF, all steps) every PASS is rock-solid CAPTURED physical evidence per §11.4.123 (§11.4.5/§11.4.69/§11.4.107) — metadata-only / config-only / absence-of-error / grep-without-runtime PASS forbidden (§11.4/§11.4.1); unclear validation method ⇒ deep-research-before-declaring-untestable per §11.4.123/§11.4.8/§11.4.99; an operator-found defect the green suite missed triggers §11.4.138. Honest boundary (§11.4.6): STEP 3 reduces the unknown-unknown surface but does not prove zero remaining defects (§11.4.118) — un-exercised cases stated as honest gaps, never silently implied clean. Composes §11.4.43/.115/.102/.130/.108/.139/.50/.85/.118/.135/.3/.123/.5/.69/.107/.138/.4/§107/§1.1. Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-146-PROPAGATION` (literal `11.4.146`) + recommended gate `CM-REPRODUCE-FIRST-THEN-EXTEND` (every closed defect's fix carries a §11.4.115 reproduce-first polarity test with captured defect-characterisation [STEP 1] + its `RED_MODE=0` GREEN confirmation [STEP 2] + an enumerated per-functionality extend case-set registered into the §11.4.135 suite [STEP 3]; a closure with the reproduce → confirm pair but NO enumerated extend case-set FAILs) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; close a defect with only the reproduce → confirm pair and no extend-case-set →

CM-REPRODUCE-FIRST-THEN-EXTEND FAILs; gate-code = separate work item).

Canonical authority: constitution submodule Constitution.md §11.4.146. Non-compliance is a release blocker. No escape hatch — no

--skip-reproduce-first , --fix-without-red , --skip-extend-to-all-cases , --reproduce-confirm-suffices ,
--defer-extend-to-release-sweep , --single-guard-suffices , --no-defect-characterisation flag.

§11.4.147 — Crashed-agent respawn-until-complete + no-work-loss registry mandate (User mandate, 2026-06-10).

Verbatim operator intent: “any agent that crashed because of something MUST BE respawned and finish its work at some point! We MUST NOT lose any work, forget about or have it corrupted!” Forensic case study (FACT): dispatched subagents died on TRANSIENT causes (API Error: Server is temporarily limiting requests rate-limit killed 5 at once, API Error: socket connection closed unexpectedly killed 2, one left PARTIAL edits — two source files written, the dependent SystemUI sliders + doc + test NOT done), each re-dispatched by pure conductor vigilance with NO mechanical guarantee — the moment conductor context is lost / compacted / the conductor itself crashes, an in-flight-but-dead work unit is silently dropped (lost / forgotten / corrupted). Every dispatched agent/subagent MUST be tracked through its full lifecycle so a crash NEVER loses, forgets, or corrupts its work; a crashed agent is NOT a completed agent (§11.4.6 — “it probably finished before dying” is a forbidden guess; abnormal termination is positive evidence the unit is OPEN). The mandate (ALL hold): **(a) durable agent REGISTRY** — every dispatched agent has an append-only machine-readable registry entry (stable id, task + §11.4.58 file-scope, declared output path(s) + §11.4.5/§11.4.69 evidence dir, dispatch timestamp, status from the CLOSED SET {dispatched | in-flight | crashed | respawned | complete}), reusing the §11.4.116 sync substrate (JSONL event stream + atomic status snapshot; “an entry with no dispatch-event cannot show complete ”); the registry is the SINGLE SOURCE OF TRUTH for “is this work owed?” and an unregistered dispatch is itself a violation; **(b) mechanical CRASH-DETECTION + RESPAWN-until-complete** — any abnormal termination (transient rate-limit/socket-close, any non-completion exit, or a terminal error after the runtime’s own retries are exhausted) flips the entry to crashed + keeps the unit OPEN, and the unit is respawned (fresh agent claims the same id + scope + preserved state) and re-respawned until an agent reaches complete ; respawn is the safe/reversible/bounded decision taken autonomously per §11.4.101 (never

blocks the loop for a human), transient causes use the project's ALREADY-DEFINED backoff budgets (never invented, §11.4.6), a deterministically-reproducible non-transient terminal error is investigated per §11.4.102 before further respawn (never a blind retry loop); **(c) PARTIAL-STATE preserve** → **§11.4.84-check** → **resume-or-clean-restart** — the crashed agent's uncommitted edits + output doc are PRESERVED (never silently discarded → lost, never blindly committed → corruption), the respawn runs the §11.4.84 quiescence check on the preserved tree (every modified file accounted-for vs scope, no mutation/ // always pass / `_mutated_*` residue, no half-written/torn artifact), then EITHER RESUMES idempotently when it passes OR CLEAN-RESTARTS from a known-good base (§9.2 pre-op backup, reversible §11.4.101) when inconsistent — nothing lost, nothing corrupted; per-agent `git worktree` isolation (§11.4.58 L4 / §11.4.84) keeps the partial tree from contaminating other streams; **(d) “crash ≠ done”**

COMPLETION criterion — a unit is DONE only when an agent reaches `complete` with its required captured evidence/output landed (§11.4.5/§11.4.69 + the §11.4.116 verdict-carries-evidence-path rule); the endless-loop done-condition (§11.4.87/§11.4.94/§11.4.97/§11.4.126) MUST NOT read satisfied while any entry is `dispatched / in-flight / crashed / respawned -not-yet- complete`, the zero-idle survey (§11.4.94) treats every non-`complete` entry as an OPEN item to reclaim, and a registry showing `complete` without landed evidence is a §11.4 PASS-bluff at the agent-lifecycle layer. Honest boundary (§11.4.6): the registry + respawn guarantee work is not lost/corrupted, NOT that it is correct — the respawned output still crosses §11.4.108 + §11.4.125/§11.4.142 review + §11.4.40 retest; §11.4.147 is the durability layer beneath those, the agent-side analogue of §11.4.144 (same detect → wait/backoff → re-attach/respawn → escalate shape). Composes §11.4.6/.58/.84/.87/.94/.97/.101/.116/.102/.108/.125/.142/.40/.128/.144/ §9.2/§1.1. Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-147-PROPAGATION` (literal `11.4.147`) + recommended gate `CM-CRASHED-AGENT-RESPAWN-TRACKED` + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; mark a `crashed` entry as the loop's done-condition `complete` without landed evidence, OR drop a dispatched agent without a registry entry → `CM-CRASHED-AGENT-RESPAWN-TRACKED` FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.147. Non-compliance is a release blocker. No escape

hatch — no --skip-agent-registry , --crash-equals-done , --no-respawn ,
--discard-partial-state , --blind-commit-partial , --forget-dead-agent , --loop-done-with-crashed-entries flag.

§11.4.148 — Workable-item integrity (status+type+id) + comprehensive structured description + bidirectional external-tracker sync + BLOCKED unblock-choices mandate (User mandate, 2026-06-10). BINDS + STRENGTHENS the workable-item discipline (§11.4.15 status / §11.4.16 type / §11.4.54 id / §11.4.21 Operator-blocked / §11.4.91 description-clarity / §11.4.93 + §11.4.95 SQLite SSoT / §11.4.106 docs_chain) into ONE integrity contract spanning DB ↔ docs ↔ external tracker, adding the operator's three emphases: **(D1) no item without a valid status + valid type + stable id — on ALL three surfaces** (an item missing any of the three in the DB, the rendered docs, OR the external tracker FAILs the validator — release-blocker; the id is the cross-surface binding key); **(D2) comprehensive structured description per item** — WHAT it is (§11.4.91 ≥6-word/≥40-char clear meaning) + HOW it manifests + HOW to reproduce (§11.4.115/.146) + ACCEPTANCE CRITERIA (§11.4.123/.69 captured-evidence verdict), in the §11.4.93 DB description column AND the docs AND the tracker, never a stub/§11.4.91-anti-pattern fragment; **(D3) BLOCKED items carry WHY + enumerated unblock CHOICES** — tightens §11.4.21: every Operator-blocked item (incl. Blocked / BLOCKED documented alias normalised to canonical Operator-blocked , never a silent fork §11.4.6) MUST enumerate the closed list of decisions/actions that would unblock it ([A]...·[B]...·[C]... , mirroring §11.4.66's 2–4-option shape); a BLOCKED item with no enumerated choices FAILs; **(D4) regular never-missed bidirectional DB ↔ docs ↔ tracker sync** — the §11.4.93/.95 git-tracked SQLite DB is the SSoT, docs + tracker DERIVED; full sync (db-to-md / md-to-db byte-identical round-trip §11.4.93 + tracker push) runs regularly + on every change, §11.4.86 drift-proof fingerprint (sha256 of the sorted item keyset, NOT mtime) gates freshness, §11.4.106 docs_chain-bound so drift is mechanically caught not vigilance-dependent; **(D5) generic external-tracker push** carries statuses (collapsed onto the tracker's native set per §11.4.33/.112 when fixed, precise value preserved in a header, never lost), types, assignee (§11.4.104 participant handle, UNSET defaults from a project env var, never hardcoded/logged §11.4.10), and sub-tasks, IDEMPOTENT (dry-run-then-real, match-by-stable-key [<ID>] prefix / id custom-field, present⇒UPDATE/absent⇒CREATE, rate-limited, credential-redacted, sink-side created=N updated=M failed=0 proof §11.4.69), the MACHINERY project-agnostic §11.4.28

(consumer registers its tracker/list/field-map at runtime). Anti-bluff §11.4: every sync + validator pass carries captured evidence §11.4.5/.69; honest boundary §11.4.6 — guarantees well-formed items + agreeing surfaces, NOT that the underlying work is correct (still crosses §11.4.108/.40/.123). Composes

§11.4.15/.16/.21/.33/.34/.54/.66/.86/.91/.93/.95/.104/.106/.112/.123/.10/.28/.69/.6/

§1.1. Classification: universal (§11.4.17) — consumer supplies its DB path, id prefix, tracker service + list/board id + field map + default-assignee env var, docs_chain context per §11.4.35. Propagation gate CM-COVENANT-114-148-PROPAGATION (literal 11.4.148) + recommended gates CM-ITEM-INTEGRITY-STATUS-TYPE-ID / CM-ITEM-COMPREHENSIVE-DESCRIPTION / CM-BLOCKED-UNBLOCK-CHOICES / CM-TRACKER-SYNC-IDEMPOTENT + paired §1.1 mutation (gate-code = separate work item). **Canonical authority:** constitution submodule

Constitution.md §11.4.148. Non-compliance is a release blocker. No escape hatch — no --item-without-status , --item-without-type , --item-without-id , --stub-description-OK , --blocked-without-choices , --skip-tracker-sync , --one-way-sync-OK , --non-idempotent-tracker-push , --hardcode-assignee flag.

§11.4.149 — Per-workable-item testing-diary mandate (User mandate, 2026-06-10). Every workable item MUST carry an append-only TESTING DIARY — a chronological record of every test run against it — distinct from §11.4.93

item_history (lifecycle STATE transitions) + §11.4.55 Reopens (reopen cycles): the diary records TEST EXECUTIONS (which may or may not change status). The mandate (ALL hold): **(a)** a test_diary table additive to the §11.4.93/.95 SQLite SSoT, one row per run soft-keyed to the item id, capturing ISO-8601-UTC

date_time + tested_by (closed set User|Operator|AI-agent|HelixQA) + result (§11.4.45 PASS|FAIL|SKIP + detail) + in-depth observations (long markdown, facts or explicit UNCONFIRMED: / PENDING_FORENSICS: §11.4.6, the background a future fix needs) + action_taken + status_changed +from/to (did this run change status + WHY) + §11.4.69 evidence_path + feature_class , with a SCHEMA CONSTRAINT making a PASS row WITHOUT a non-empty evidence path impossible (a PASS-bluff rejected by the schema itself); **(b)** BOTH an in-depth diary doc + a derived at-a-glance summary VIEW (total/pass/fail/skip + last-verdict + last-run + status-changes + distinct testers/feature-classes, DERIVED never duplicated §11.4.93) feeding §11.4.132 risk-ordering; **(c)** four-format per-item Diary / Diary_Summary exports §11.4.65 + §11.4.86-drift-proof-fingerprinted (sha256 of the sorted diary keyset) + §11.4.106 docs_chain-bound so a diary row

not exported/pushed FAILs a gate (never-missed); **(d)** external-tracker SUB-TASK model — the item task's description stays the item (§11.4.148 D2, NOT polluted with diary text), each run a CHILD sub-task (type `task`) with a `{TODO|In-progress|Completed}` lifecycle (collapsed per §11.4.33/.112), observations + action + an Evidence: line (path not raw artefact §11.4.10/.13) + a diary-entry idempotency key, reusing the §11.4.148 D5 idempotent rate-limited credential-redacted sink-side-proven push, mapper project-agnostic §11.4.28; **(e)** MINIMAL-LLM deterministic bash/Go tooling (zero LLM in the data path — the `observations` prose is authored by whoever ran the test; tooling only stores/renders/pushes/validates), 100% test-covered §11.4.27 (unit + integration-against-the-real-tracker with an honest §11.4.3 SKIP-when-token-absent never a faked PASS + export + HelixQA Challenge + paired §1.1 mutation) so a diary PASS-bluff is mechanically impossible. Honest boundary §11.4.6 — the diary guarantees a complete auditable per-item test-history, NOT that any single run's verdict is correct (rests on its own §11.4.69/.107/.123 evidence). Composes §11.4.6/.27/.45/.50/.55/.65/.69/.86/.93/.95/.106/.107/.123/.132/.148/.10/.13/.28/.33/.112/ §1.1. Classification: universal (§11.4.17) — consumer supplies its DB path, diary doc layout, export formats, tracker sub-task field map, docs_chain context per §11.4.35. Propagation gate `CM-COVENANT-114-149-PROPAGATION` (literal `11.4.149`) + recommended gates `CM-TEST-DIARY-SYNC` / `CM-DIARY-PASS-REQUIRES-EVIDENCE` + paired §1.1 mutation (gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.149. Non-compliance is a release blocker. No escape hatch — no `--skip-testing-diary`, `--diary-without-evidence`, `--no-diary-summary`, `--diary-without-tracker-subtask`, `--llm-in-diary-data-path`, `--diary-export-optional` flag.

§11.4.150 — Mandatory deep multi-angle web research per change/issue, before declaring fixed or structural (User mandate, 2026-06-11). Verbatim operator mandate: “For every single issue we fix or improvement we make — besides tight systematic-debugging, fixing, code review by independent agents, comprehensive tests — we MUST ALWAYS do deep web research from various angles! No matter how big/small/simple/complex, dig deep the internet for articles, technical documentation, APIs, open-source code — EVERYTHING that can help make the best possible solution OR confirm we don't have a more serious problem we're unaware of! We MUST do everything possible to STOP the constant issue-reopening and finally start closing items as fixed+working! Do this ALWAYS in

parallel with the main work stream!" For EVERY fix / improvement / change / closure — no matter how big, small, simple, or complex — the agent MUST, IN ADDITION to tight systematic-debugging (§11.4.102), the fix, independent-agent code review (§11.4.125 / §11.4.142), and comprehensive multi-layer tests (§11.4.4(b) / §11.4.40), perform a DOCUMENTED **deep multi-angle web research pass** digging the internet from VARIOUS angles — articles, official + vendor technical documentation, API references, standards, issue trackers, maintainer guidance, reusable open-source code — to BOTH (i) discover the best possible solution AND (ii) confirm there is NOT a more serious underlying problem the team is unaware of. Forensic case study (FACT, 2026-06-11): a subtitle-on-secondary-display goal verdicted `Won't-fix: structurally-impossible` (§11.4.112 secure-surface pixel-blanking) was OVERTURNED by a deep multi-angle research pass that surfaced the real OCR-of-the-PRIMARY-display path — a “we can’t” became a shipped capability; the canonical anti-pattern of a structural verdict reached for want of research. The mandate (ALL hold): **(A) No closure-as-fixed/structural WITHOUT a documented deep-research pass** — no item may be marked `Fixed / Implemented / Completed` (§11.4.33) OR classified `structurally-impossible won't-fix` (§11.4.112) until a deep multi-angle pass is documented; §11.4.150 makes §11.4.8's pass UNCONDITIONAL (every item, however trivial, at the closure AND structural-verdict gates). **(B) Multiple angles, not a single lookup** — ≥ 2 genuinely-distinct angles (official-docs / standards / known-bug-trackers / alternative-approach / failure-mode / security / performance / platform-constraint / open-source-precedent) so a single-source confirmation-bias miss (§11.4.145) cannot pass; a one-link drive-by is NOT deep multi-angle. **(C) Confirm-no-bigger-problem, not just find-a-fix** — explicitly seek evidence the fix does not mask a deeper defect AND the closure/verdict is genuinely safe; “we found nothing worse” requires the enumerated search (§11.4.118). **(D) Latest-source + cited** — LATEST authoritative versions per §11.4.99 (never training-data/memory), each cited by URL + access date in the research artefact AND the closure commit footer (`Deep-research <date>: <urls>` OR the literal `NO external solution found – original work` per §11.4.8). **(E) Reopen-breaking is the PURPOSE** — STOP the constant Fixed ↔ Reopened churn (§11.4.34 / §11.4.55); a reopen whose root cause a deep pass would have surfaced is a §11.4.150 miss; composes §11.4.7 (demotion-evidence) + §11.4.112 (structural verdict now REQUIRES the cited-authorities pass). **(F) ALWAYS in parallel with the main stream** — background subagent-driven (§11.4.70 / §11.4.20 / §11.4.103 / §11.4.89) concurrent with main fix/build/test work, NEVER serialising it or stalling

the loop (§11.4.94 / §11.4.97 / §11.4.101 / §11.4.126). **(G) Apply ASAP to every workable item** — retroactively + going forward across the §11.4.93 / §11.4.95 SSoT; items closed without a documented pass are re-audited in the §11.4.40 / §11.4.42 release-gate sweep. Honest boundary (§11.4.6): the pass reduces the unknown-unknown surface + breaks the most common reopen causes — it does NOT prove zero remaining defects (§11.4.118) and does NOT replace §11.4.108 runtime-signature verification, §11.4.125 / §11.4.142 review, or §11.4.40 retest; it is the research layer every fix/closure/structural-verdict additionally crosses, and “we probably don’t have a bigger problem” without the enumerated multi-angle search is a guess (§11.4.6), never a finding. Composes §11.4.8 / §11.4.99 / §11.4.123 / §11.4.118 / §11.4.145 / §11.4.125 / §11.4.142 / §11.4.7 / §11.4.112 / §11.4.34 / §11.4.55 / §11.4.70 / §11.4.20 / §11.4.89 / §11.4.103 / §11.4.40 / §11.4.42 / §11.4.93 / §11.4.95 / §11.4.6 / §1.1. Classification: universal (§11.4.17) — the consuming project supplies its concrete research corpora, angle set, item tracker, and closure-commit-footer convention per §11.4.35. Propagation gate `CM-COVENANT-114-150-PROPAGATION` (literal `11.4.150`) + recommended gate `CM-DEEP-RESEARCH-PER-ISSUE` (every closed/structural-verdicted item carries a documented multi-angle deep-research artefact + cited-source closure footer, run in parallel, before the closure/structural verdict is accepted) + paired §1.1 mutation (strip the literal → propagation gate FAILs; close an item or reach a `structural` verdict with no documented multi-angle research artefact / cited footer → `CM-DEEP-RESEARCH-PER-ISSUE` FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.150. Non-compliance is a release blocker. No escape hatch — no `--skip-deep-research`, `--single-source-suffices`, `--trivial-change-no-research`, `--close-without-research`, `--structural-without-research`, `--serialise-research`, `--research-from-memory-OK` flag.

§11.4.151 — Project-prefixed release-tag/version-naming mandate (User mandate, 2026-06-12). Verbatim operator mandate: “Every release tag and version name we create — on the main repository and on every Submodule we own — MUST be prefixed with the project’s release prefix, e.g. `myproject-1.0.0-dev-0.0.1`. The prefix MUST come from `HELIX_RELEASE_PREFIX` in our `.env` if it is set, otherwise from the lowercased project root directory name. The SAME prefix MUST be used across the main repo and all owned Submodules in one release so a release is greppable across every repository.” Every release tag AND every version name created on the main repository AND on every owned-by-us

submodule (§11.4.28) MUST be prefixed with the project's release prefix, form `<PREFIX>-<version>` (canonical example: `myproject-1.0.0-dev-0.0.1`), so a release is identifiable + greppable across every repository it spans (one `git tag -l '<PREFIX>-*'` enumerates the whole release surface). **Prefix resolution order (closed-set, deterministic — §11.4.6 no-guessing):** (1) `HELIX_RELEASE_PREFIX` from the project's `.env` — authoritative when set; `.env` is git-ignored per §11.4.30 and the variable is documented in the tracked `.env.example` (a §11.4.77 re-obtain mechanism, never committed); (2) fallback = the lowercased snake_case form of the project root directory name (no spaces) per §11.4.29 — used whenever the env var is unset/empty, so a prefix is ALWAYS resolvable from the checkout with zero operator input. **The prefix MUST be IDENTICAL across the main repo and all owned submodules within a single release** — a release that tags the main repo `<PREFIX>-1.2.0` while tagging an owned submodule with a different/unprefixed value is a §11.4.151 violation (the cross-repo grep no longer enumerates the release). Version codes increment monotonically within the prefix (`<PREFIX>-...-0.0.1` → `<PREFIX>-...-0.0.2` → ...), never reset, never skipped. Honest boundary (§11.4.6): the prefix guarantees a release is identifiable + uniform across every repository, NOT that its contents are correct — the tag is still created only after the §11.4.40 full-suite retest GREEN and reaches every upstream via the §11.4.113 merge-onto-latest-main path (NEVER a force-push), fanned out per §2.1. Composes §2.1 / §11.4.29 / §11.4.30 / §11.4.40 / §11.4.113 / §11.4.126 (the release-scope terminal condition is a published, prefixed tag). Classification: universal (§11.4.17) — the consuming project supplies its concrete prefix value + the `HELIX_RELEASE_PREFIX` env var per §11.4.35. Propagation gate `CM-COVENANT-114-151-PROPAGATION` (literal `11.4.151`) + recommended gate `CM-RELEASE-PREFIX-NAMING` (every release tag/version on the main repo + every owned submodule carries the resolved `<PREFIX>-` prefix, identical across the release) + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; create an unprefixed or differing-prefix release tag → `CM-RELEASE-PREFIX-NAMING` FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.151. Non-compliance is a release blocker. No escape hatch — no `--no-release-prefix`, `--unprefixed-tag`, `--prefix-optional`, `--differing-submodule-prefix` flag.

§11.4.152 — Crashlytics-recorded-data continuous monitoring + systematic-debug + regression-test-coverage mandate (User mandate, 2026-06-13).

Verbatim operator intent: “For every project that has Firebase Crashlytics enabled /

wired, we MUST continuously monitor ALL of the Crashlytics-recorded data — crashes, ANRs, performance traces, and non-fatals — systematically debug each, fix and improve, and cover everything with validation and verification tests. This MUST be checked regularly, with no false results and no bluff of any kind!” Every project that has Firebase Crashlytics enabled/wired (SDK linked into a shipping artifact, crash+non-fatal+ANR reporting active) MUST treat the Crashlytics console as a first-class captured-evidence channel from real end-user devices and continuously process every datum it records — an open Crashlytics issue with a green test suite is a §11.4 PASS-bluff at the field-telemetry layer (a real user hitting a broken feature while the suite reports green). STRENGTHENS+COMPLETES §11.4.47: §11.4.47 owns the periodic REVIEW + dedup + Issue-creation pass that SURFACES Crashlytics/Analytics/Performance findings; §11.4.152 owns what happens to each surfaced item AFTER — the systematic-debug → fix/improve → regression-test-coverage lifecycle that drives it to a proven regression-immune closure (§11.4.47 finds it; §11.4.152 fixes it + proves it stays fixed; both mandatory, neither substitutes). The mandate (ALL hold): (1) **continuous monitoring of ALL four surfaces** — fatal crashes, ANRs, performance traces/regressions, AND non-fatals (skipping any one is a PASS-bluff; non-fatals are the silent class — every catch/fallback/recovered-error path is a real degraded UX and MUST be triaged, not ignored because the app did not crash); (2) **systematic-debugging of each (reproduce-before-fix)** per §11.4.102 Iron Law + §11.4.115 RED-baseline-on-the-broken-artifact (the recorded stacktrace/ANR-trace/non-fatal-context IS the §11.4.5/.69 captured defect-present evidence) BEFORE the fix — a fix without a prior reproducing test is a §11.4.43/.123 violation; unreproducible ⇒ deep-research-before-declaring-untestable §11.4.123/.150; (3) **a fix/improvement for every confirmed issue** against its proven root cause (§11.4.9/.43/.108), “acknowledged in console” is not a fix, muting a recurring issue without a tracked rationale is a §11.4.6/.90 violation; (4) **validation+verification regression-test coverage per closed issue** — in the SAME commit as the fix register a permanent §11.4.135 regression guard (§11.4.115 polarity test: RED_MODE=1 captures the recorded defect on the pre-fix artifact, RED_MODE=0 standing GREEN guard asserting ABSENT) exercising the same user-reachable path the stacktrace identifies, with rock-solid captured evidence §11.4.5/.69/.107/.123 + a closure log citing the console issue id/URL + root-cause + fix commit + the validation/verification test paths; a Crashlytics issue marked resolved WITHOUT a falsifiable regression test is FORBIDDEN (the silent-recurrence vector, §11.4.138 operator-escape class applied to field telemetry); (5) **regular cadence** reusing the §11.4.47 five-trigger set

(pre-build/pre-flash/pre-distribute/pre-tag blocking; daily + post-deployment-burn-in non-blocking) — a one-time sweep never repeated is a violation; (6) **no false results, no bluff** (§11.4.1/6 — “no new issues” requires the enumerated monitored-surfaces+window §11.4.118; “fixed” requires the RED → GREEN flip §11.4.115/.130; a guard whose §1.1 mutation does not FAIL it is a bluff gate). Honest boundary §11.4.6: processing every recorded issue breaks the silent-recurrence vector — it does NOT prove zero remaining field defects nor replace §11.4.108/.40; an issue is “closed” only when its guard is GREEN on a clean artifact, never when the console mark is flipped. Classification: universal (§11.4.17) — the consuming project supplies its Crashlytics handle, console-access credential (§11.4.10, never logged), severity table, and per-issue closure-log path per §11.4.35; the reference project-level instantiation is a consuming project’s §6.O (Crashlytics-Resolved Issue Coverage — per-issue validation test + Challenge test + `.lava-ci-evidence/crashlytics-resolved/<date>-<slug>.md` closure log) + §6.AC (Comprehensive Non-Fatal Telemetry — every catch/fallback records a non-fatal with triage context). Composes §11.4.1/.5/.6/.9/.34/.40/.43/.47/.69/.90/.102/.107/.108/.115/.118/.123/.130/.135/.138/.150/ §1.1. Propagation gate `CM-COVENANT-114-152-PROPAGATION` (literal `11.4.152`) + recommended gate `CM-CRASHLYTICS-ISSUE-FULLY-COVERED` (every closed Crashlytics issue carries a systematic-debug root-cause record + a registered §11.4.135 regression guard + a closure-log entry citing the console issue id/URL + the validation/verification test paths; a console-resolved issue with no falsifiable regression guard FAILs) + paired §1.1 meta-test mutation (gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.152. Non-compliance is a release blocker. No escape hatch — no `--skip-crashlytics-monitoring`, `--monitor-crashes-only`, `--skip-non-fatals`, `--resolve-without-regression-test`, `--console-mark-is-fixed`, `--monitor-once`, `--mute-without-rationale` flag.

§11.4.153 — Comprehensive per-feature Status + Status_Summary document set with mandatory video-recording confirmation (User mandate, 2026-06-15).

Every project MUST maintain, under `docs/features/`, a comprehensive feature Status document set (`Status.md` + its §11.4.56 `Status_Summary.md` companion) enumerating EVERY system component, EVERY client app/binary/surface (TUI / CLI / Web / desktop / mobile / API / gRPC / library / submodule / infrastructure), and EVERY feature — including features ported from any incorporated CLI-agent / submodule catalogue (§11.4.74) — with NO single feature

left out. (1) **Total categorized coverage** — per-feature table organized Component → Category, reconciled against the actual codebase (a code-present feature missing from the table, or a row with no code, is a §11.4.6/§11.4.118 bluff).

(2) **Per-feature fields** — Component / Feature / Category / Implementation / Wiring (genuinely end-user-reachable, not merely compiled, §11.4.108) / Real-use / Tests-coverage (four-layer §11.4.4(b)) / Validation (PASS / FAIL / SKIP / PENDING_FORENSICS / OPERATOR-BLOCKED per §11.4.45) / **Video-recording confirmation** (path to the real-use video or honest gap marker). (3) **Mandatory per-feature real-use video** — every user-visible “confirmed” claim backed by a recorded real-use video of a genuine end-user scenario (real prompts → real LLM/ service responses → real results), NEVER a frozen/stale frame (§11.4.107), NEVER faked/mock/demo-loop/bluff response or LLM error passed off as success (§11.4.2/§11.4.5), stored at the project-declared recording path (§11.4.35); a confirmed row with no real video or a bluff video = §11.4 PASS-bluff; autonomous-infeasible ⇒ honest §11.4.3/§11.4.52 SKIP + tracked migration item, NEVER a faked confirmation. (4) **Video-analysis remediation loop** — every video analysed; any defect it surfaces triggers §11.4.102 systematic-debug → fix → §11.4.146 retest → re-record → clean GO per §11.4.134 before “confirmed” (a video exposing a broken feature is a §11.4.4 test-interrupt). (5) **Always-in-sync** — §11.4.45-class roster/corpus-backed, §11.4.106 docs_chain-bound + §11.4.86 drift-proof fingerprint (sha256 of sorted feature-key roster AND sorted video-artefact roster, NOT mtime), re-syncs out-of-the-box; stale = violation. (6) **Four-format export** — HTML + PDF + **DOCX** (this doc class ADDS DOCX to the §11.4.65 HTML+PDF set; other classes unchanged), in sync per §11.4.60. (7) Follows §11.4.44/.45/.56/.57/.59/.60. (8) **MP4 format REQUIRED**. All video confirmations MUST be in .mp4 format (H.264). Window-specific capture ONLY (§11.4.159(A)). Vision validation REQUIRED (§11.4.159(D)). .cast files are supplementary only. Honest boundary §11.4.6 — guarantees a complete video-confirmed always-synced ledger, NOT per-feature correctness (rests on each row’s §11.4.69/.107/.123 evidence) and NOT a §11.4.40 retest substitute. Classification: universal (§11.4.17). Composes §11.4.2/.5/.44/.45/.52/.56/.57/.59/.60/.65/.86/.102/.106/.107/.108/.118/.123/.134/.146/ §1.1. Propagation gate CM-COVENANT-114-153-PROPAGATION (literal 11.4.153) + recommended gates CM-FEATURE-STATUS-COMPLETE + CM-FEATURE-STATUS-VIDEO-CONFIRMED + paired §1.1 mutation.

Canonical authority: constitution submodule `Constitution.md` §11.4.153. Non-compliance is a release blocker. No escape hatch — no `--skip-feature-status`, `--feature-without-video`, `--frozen-video-OK`, `--bluff-response-video-OK`, `--skip-video-analysis`, `--feature-ledger-incomplete-OK`, `--no-docx-export`, `--allow-feature-ledger-drift` flag.

§11.4.154 — Window-scoped capture + fresh-corpus rotation for feature/QA recordings (User mandate, 2026-06-15). Verbatim: “recording you perform does record the window containing apps and services, not the whole desktop or monitor screen!” + “All old recording files MUST BE removed when new one starts!” Refines §11.4.2/.5/.107/.153 recording discipline with two capture-hygiene invariants. **(A) Window-scoped, NOT whole-screen** — every feature/QA video MUST capture ONLY the window/surface of the app/service under test (GUI window / CLI-TUI terminal pane / web tab-viewport / device-emulator-simulator frame), NEVER the whole desktop/monitor or unrelated windows; whole-desktop capture leaks operator-private content (§11.4.10/.83), dilutes the §11.4.107 liveness/freeze oracle, and breaks the §11.4.137 OCR/ROI content oracle. Target the window/region by stable identity (window id/title, device serial, browser context, tmux target) per §11.4.111 — never a fixed full-screen device index. Platform genuinely cannot capture below whole-screen ⇒ honest §11.4.3 SKIP + tracked migration item, never a whole-screen pass-off. **(B) Fresh-corpus rotation** — when a new recording run for a scope begins, the agent’s OWN prior in-scope stale recordings at the raw recording path MUST be removed FIRST so the live corpus reflects the current run (§11.4.107 not-stale + §11.4.86 roster-freshness). Honest boundary (§11.4.6 + §9.2): “remove old” = the agent’s own prior recordings for the SAME scope/project ONLY — NEVER another project’s/scope’s/operator-authored files; uncertain ⇒ surface, don’t delete (§11.4.122); committed `docs/qa/<run-id>/` evidence (§11.4.83) is the durable record, NOT rotated. Classification: universal (§11.4.17). Composes §11.4.2/.5/.10/.83/.86/.107/.111/.122/.128/.137/.153/§9.2/.6. Propagation gate `CM-COVENANT-114-154-PROPAGATION` (literal `11.4.154`) + recommended gate `CM-WINDOW-SCOPED-FRESH-CORPUS-RECORDING` + paired §1.1 mutation.

Canonical authority: constitution submodule `Constitution.md` §11.4.154. Non-compliance is a release blocker. No escape hatch — no `--whole-screen-capture-OK`, `--skip-window-scope`, `--keep-stale-recordings`, `--no-corpus-rotation`, `--full-desktop-recording` flag. **(C) MP4 auto-conversion**

REQUIRED. Any `.cast` file produced MUST be auto-converted to `.mp4` via `agg + ffmpeg` immediately after capture. The `.mp4` is the primary evidence; `.cast` is supplementary only.

§11.4.155 — Project-name-prefixed feature/QA recording filenames (User mandate, 2026-06-15). Verbatim: “All recorded videos MUST START with prefix: the PROJECT NAME (ALWAYS USE THE PROJECT NAME). Project name MUST be obtained according to the constitution's own project-name resolution.” Every recorded video the project produces — every feature/QA real-use recording (§11.4.153), every window-scoped capture (§11.4.154), every always-on device recording (§11.4.128), every raw or curated artefact at the project-declared recording path (§11.4.35) + the committed `docs/qa/<run-id>/` trail (§11.4.83) — MUST have a filename that STARTS WITH the PROJECT-NAME prefix, ALWAYS; an unprefixed recording is a §11.4.155 violation (a multi-project/scope corpus on one host per §11.4.128/.103 becomes un-greppable + un-attributable — the §11.4.151 identify-and-grep failure on the recording-corpus axis). **Prefix resolution (closed-set, deterministic — §11.4.6, IDENTICAL to §11.4.151):** (1)

`HELIX_RELEASE_PREFIX` from `.env` (authoritative, git-ignored §11.4.30, documented in tracked `.env.example` §11.4.77) else (2) lowercased snake_case project-root dir name §11.4.29 — ALWAYS resolvable, zero operator input. SAME prefix for EVERY recording in a checkout so `ls '<PREFIX>--- '*` enumerates the corpus; canonical form `<PREFIX>---<feature-or-scope>---<run-id>.<ext>` (`---` delimits the prefix unambiguously); MUST equal the §11.4.151-resolved release-tag prefix (divergence is itself a §11.4.155 violation — one project, one name). Honest boundary (§11.4.6): the prefix guarantees attribution + greppability, NOT content validity (still §11.4.107/.137/.153) and does NOT relax §11.4.154's window-scope/rotation (rotation removes the agent's OWN `<PREFIX>--- *` only; foreign/operator files surfaced not deleted §11.4.122/§9.2). Classification: universal (§11.4.17). Composes §11.4.151/.128/.153/.154/.111/.83/.6/.29/.30/.35/.77/.86/§1.1. Propagation gate `CM-COVENANT-114-155-PROPAGATION` (literal `11.4.155`) + recommended gate `CM-RECORDING-PROJECT-NAME-PREFIX` + paired §1.1 mutation.

Canonical authority: constitution submodule `Constitution.md` §11.4.155. Non-compliance is a release blocker. No escape hatch — no `--no-recording-prefix`, `--recording-without-project-name`, `--unprefixed-recording`, `--prefix-optional-for-recording`, `--differing-recording-prefix` flag.

§11.4.156 — All CI/CD automation (GitHub Actions / GitLab pipelines / equivalents) MUST be disabled (User mandate, 2026-06-15). Verbatim operator mandate: “Any GitHub actions or GitLab pipelines MUST BE disabled! Add this critical mandatory rule / mandatory constraint into the root constitution Submodule, commit and fetch all its changes to all upstreams and make sure we respect and follow this rule (we do apply it) ASAP!!!” Every repository this Constitution governs — main repo, this constitution submodule, every owned + nested submodule we author and push — MUST ship with ALL server-side CI/CD automation DISABLED: no push to any owned upstream may trigger a GitHub Actions run, GitLab pipeline, or equivalent (Jenkins/CircleCI/Travis/Drone/Woodpecker/Bitbucket/Azure, any on: push / schedule / workflow_dispatch). GENERALISES + makes ABSOLUTE the §11.4.75 Layer-5 posture (remote CI DISABLED, workflow preserved at a ...disabled-local-only non- .yaml name a provider ignores) across ALL governed repos; enforcement migrates to the LOCAL §11.4.75 git-hook ritual + §11.4.40 pre-tag sweep, never a remote runner. ALL hold: **(A)** zero active .github/workflows/*.yaml|yml / .gitlab-ci.yml / .gitlab/** / equivalent at the ROOT of any governed repo/submodule (the only place a provider executes); **(B)** “disabled” = a push triggers ZERO runs — delete OR rename to a non-trigger name (§11.4.75 .disabled / .disabled-local-only); a live- on: + if:false workflow still queues runs, NOT compliant; **(C)** scope = repos we author+push — vendored/third-party nested configs below the root (AOSP external/** , prebuilts/** , vendored submodules) are INERT (a provider never runs a non-root config), OUT of scope, MUST NOT be mass-edited (§11.4.29 vendor-exempt); test = “does a push to OUR upstream trigger a run?” yes⇒disable, inert⇒document+leave (§11.4.6 verify-not-assume); **(D)** no new CI may be added (release blocker); **(E)** pre-push verify git ls-files | grep -E '^\.github/workflows/.*\.ya?ml\$|^\.gitlab-ci\.yaml\$' empty for authored repos, §11.4.109-class PreToolUse guard + gate enforce mechanically. Honest boundary (§11.4.6): file-level disabling stops FILE-triggered runs, NOT provider-side server settings (org-default required workflows, branch-protection required checks, provider scheduled exports) — the operator turns those off; the agent documents what it cannot reach, never claims unachieved completeness. Composes §11.4.75/.29/.6/.40/.42/.109/.113/§2.1. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-156-PROPAGATION (literal 11.4.156) + recommended gate CM-NO-ACTIVE-CI + paired §1.1 meta-test mutation (strip the literal → propagation gate FAILs; add a root .github/workflows/x.yaml to an

authored repo → `CM-NO-ACTIVE-CI FAILS`). **Canonical authority:** constitution submodule `Constitution.md` §11.4.156. Non-compliance is a release blocker. No escape hatch — no `--allow-ci` , `--enable-workflow` , `--keep-pipeline` , `--remote-ci-OK` , `--ci-exempt` flag.

§11.4.157 — GEMINI.md maintained in lockstep with CLAUDE.md /

AGENTS.md / QWEN.md (User mandate, 2026-06-15).

Verbatim operator mandate: “Make sure with CLAUDE.md, AGENTS.md, QWEN.md we maintain GEMINI.md too! Add this mandatory fact / rule to root constitution Submodule we are inheriting / extending - CONSTITUTION.md, CLAUDE.md, QWEN.md, AGENTS.md, GEMINI.md and other related relevant files!” Forensic FACT (2026-06-15): when §11.4.156/§11.4.157 were authored, Constitution/CLAUDE/AGENTS/QWEN carried the family through §11.4.155 but GEMINI.md had silently drifted to §11.4.141 — 14 mandates (§11.4.142–155) never propagated — a §11.4 propagation-bluff (a Gemini-CLI agent reads a stale Constitution). GEMINI.md is a FIRST-CLASS governance context carrier EQUAL to CLAUDE.md/AGENTS.md/QWEN.md, never optional/best-effort. ALL hold: **(A)** five-carrier lockstep — no governance change is complete until GEMINI.md carries it alongside the other three mirrors; GEMINI.md is added to the §11.4.26 propagation + cross-reference set explicitly; **(B)** no silent drift — GEMINI.md lagging the other mirrors’ highest rule is a §11.4.157 violation (§11.4.65-class), back-fill required; **(C)** equal status — GEMINI.md restates the SAME literal `11.4.N` anchors the propagation gates require (§11.4.35), fleet count INCLUDES GEMINI.md; **(D)** consumer projects’ own CLAUDE/AGENTS/QWEN/GEMINI bind too (§11.4.35). Honest boundary (§11.4.6): the §11.4.142–155 GEMINI.md back-fill is a tracked release-blocking remediation; claiming GEMINI.md “in sync” while the back-fill is incomplete is itself a §11.4.157 violation. Composes §11.4.26/.35/.17/.44/.65/.140/.156/§1.1. Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-157-PROPAGATION` (literal `11.4.157` , GEMINI.md INCLUDED) + recommended gate `CM-GEMINI-MD-LOCKSTEP` + paired §1.1 meta-test mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.157. Non-compliance is a release blocker. No escape hatch — no `--skip-gemini-md` , `--gemini-optional` , `--gemini-lag-OK` , `--four-carrier-suffices` flag.

§11.4.158 — Intensive all-feature/flow/edge-case video-recording + read-the-screen content-verification mandate (User mandate, 2026-06-16). Every project MUST be covered by intensive automated testing that exercises + RECORDS every feature/flow/use-case/edge-case (valid/invalid/boundary/concurrent/failure-

injection §11.4.85), each recording showing the feature GENUINELY WORKING with REAL results (real prompts → real responses → real outputs §11.4.153) and NO false/simulated/stale/frozen result (§11.4.2/.5/.107) — a recording showing an error/bluff/non-working feature is a FINDING (→ §11.4.153(4)/§11.4.4 fix → retest → re-record), never a confirmation. The testing System MUST ACTUALLY READ every shown log line / message / UI label / dialog / toast / status text and VERIFY it is a genuine working result (OCR/ROI §11.4.117/.137 + confidence floor for pixel surfaces; direct capture for terminal/log; §11.4.107(10) self-validated golden-good/golden-bad analyzer) — “a video was produced” is NOT evidence, “the System read the screen + confirmed a genuine result” is. HelixQA (§11.4.27) MUST drive this exercise → record → read → score pass, PASS only on a read-confirmed genuine result + captured artefact path (§11.4.69). Default recording save path = `$HOME/Downloads` (host user’s home Downloads, resolved at runtime never hardcoded) unless a project declares an override per §11.4.35; §11.4.155 project-prefix + §11.4.154 window-scope/rotation + §11.4.128 git-ignored raw corpus + §11.4.83 curated docs/qa evidence apply. **Vision analysis MANDATORY for every recording** — after every recording, the agent MUST read the terminal output / video content and verify the feature actually works. A recording without a vision-confirmed verdict is a §11.4.158 violation. Composes

§11.4.2/.5/.25/.27/.52/.69/.83/.85/.107/.108/.117/.118/.128/.137/.138/.153/.154/.155/
§1.1. Classification: universal (§11.4.17). Propagation gate

`CM-COVENANT-114-158-PROPAGATION` (literal `11.4.158`) + recommended gates

`CM-INTENSIVE-RECORDING-COVERAGE` +

`CM-RECORDING-CONTENT-READ-VERIFIED` + paired §1.1 mutation. **Canonical**

authority: constitution submodule `Constitution.md` §11.4.158. Non-compliance is a release blocker. No escape hatch — no `--skip-recording-coverage`, `--video-without-content-read`, `--happy-path-recording-suffices`, `--recording-path-anywhere`, `--unread-recording-OK`, `--skip-helixqa-read` flag.

§11.4.159 — Mandatory window-specific video recording + vision validation

mandate (User mandate, 2026-06-20). Every feature test, validation, verification, challenge, and QA session that produces video evidence MUST comply with ALL of the following: **(A) Window-specific recording ONLY** — every video MUST record ONLY the target application window (Terminal pane, TUI app, browser tab, emulator frame), NEVER the whole desktop/monitor screen; use macOS `screencapture -l<window_id>`, Linux `xdotool + ffmpeg`, or equivalent; whole-desktop capture

leaks operator-private content (§11.4.10), dilutes the §11.4.107 liveness oracle, and breaks §11.4.137 OCR/ROI; platform genuinely cannot capture below whole-screen => honest §11.4.3 SKIP + tracked migration item. **(B) MP4 format REQUIRED** — `.mp4` (H.264, `movflags +faststart`, `pix_fmt yuv420p`); `.cast` files are supplementary only; auto-convert via `agg + ffmpeg` per §11.4.154(C). **(C) Project-name prefix REQUIRED** — filename MUST start with project name in snake_case from `HELIX_RELEASE_PREFIX` in `.env` (§11.4.151) or lowercased project root dir name (§11.4.29); format `<project_name>-<feature>-<timestamp>.mp4`; unprefixed = violation. **(D) Mandatory vision validation** — after EVERY recording, the agent MUST read the terminal output / video content and verify: (i) feature ACTUALLY WORKS, (ii) LLM responses are REAL, (iii) all tests show PASS, (iv) no “TODO implement” / “simulate” / “for now” patterns, (v) output demonstrates end-user working feature; MUST produce a verdict (PASS/FAIL) with evidence path (§11.4.69). **(E) Terminal window cleanup** — after each recording, dismiss/close ONLY the Terminal window used for that recording (window-specific close via `osascript` / `xdotool`); MUST NOT close windows belonging to other processes. **(F) Real results ONLY** — every recording MUST show REAL working features; errors/empty output/simulated responses trigger fix → retest → re-record. **(G) Re-runnable evidence** — the command shown MUST be re-runnable to produce the same results. **(H) Fresh-corpus rotation** — remove agent’s own prior in-scope stale recordings FIRST (§11.4.154). **(I) Content verification MANDATORY — not duration-based** — the value of a recording is NOT its duration but its CONTENT; a recording MUST demonstrate the ACTUAL feature being used with REAL results; before accepting ANY recording, the agent MUST verify: expected output patterns ARE present (e.g., test PASS lines, API response data, feature-specific output), the feature ACTUALLY WORKS as demonstrated (not just “something ran”), LLM responses are REAL content (not simulated, not placeholder, not empty), every claim of “working” is backed by visible evidence; a 5-second recording showing a feature working correctly is MORE valuable than a 60-second recording of empty terminal; duration is NOT a proxy for quality. **(J) Expected-content specification REQUIRED** — before recording, the agent MUST specify what content SHOULD appear (expected patterns, expected test results, expected API responses); after recording, the agent MUST verify these patterns ARE present; if expected content is MISSING, the recording is REJECTED regardless of duration. **(K) Content-verification recording workflow** — mandated workflow: (1) SPECIFY expected content patterns, (2) RECORD the feature execution, (3) EXTRACT all text from the recording, (4) VERIFY expected patterns

are present, (5) CHECK for simulated/placeholder content, (6) ACCEPT only if ALL patterns found AND zero bluffs detected, (7) REJECT and re-record if ANY pattern missing or bluff detected. **(L) Root cause analysis REQUIRED for rejected recordings** — when a recording is rejected (missing expected content, bluff detected, or empty capture), the agent MUST investigate WHY before re-recording per §11.4.102; determine the root cause (timing issue, wrong command, tool failure, etc.) and fix it; simply re-recording without understanding WHY the first attempt failed is a §11.4.159 violation. **(M) Real-time monitoring RECOMMENDED** — for complex features, use real-time monitoring that analyzes output DURING recording (not after); this catches issues immediately and allows corrective action before the recording completes. Classification: universal (§11.4.17). Composes §11.4.2/.3/.5/.10/.29/.69/.83/.107/.111/.128/.137/.151/.153/.154/.155/.158/§1.1. Propagation gate `CM-COVENANT-114-159-PROPAGATION` (literal `11.4.159`) + recommended gate `CM-WINDOW-VIDEO-VALIDATED` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.159. Non-compliance is a release blocker. No escape hatch — no `--whole-screen-ok`, `--cast-only`, `--skip-vision-validation`, `--no-cleanup`, `--simulated-recording-ok`, `--unprefixed-recording` flag.

§11.4.160 — Vision-verified recording + HelixQA bridge mandate (User mandate, 2026-06-21). Compact summary: every video recording for feature/QA evidence MUST be processed through a vision/OCR pipeline that reads on-screen content and confirms expected results BEFORE acceptance; the recording system MUST provide a bridge feeding captured frames to HelixQA's test infrastructure (or equivalent) for automated read-the-screen verification against SPECIFY-phase expected patterns (§11.4.159(J)); capture frames at ≤5s intervals, run OCR/vision with self-validated analyzer (§11.4.107(10)), compare text against patterns, produce per-frame PASS/FAIL with evidence path, surface failures immediately for re-record (§11.4.159(L)). Project-configured parameters per §11.4.35. Honest boundary: does NOT replace §11.4.5 quality analysis nor §11.4.108 runtime-signature; FLAG_SECURE uses §11.4.117 proxy oracle. Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-160-PROPAGATION` + recommended gate `CM-VISION-VERIFIED-RECORDING-BRIDGE` + paired §1.1 mutation.

§11.4.161 — Rootless container runtime mandate (User mandate, 2026-06-21). Compact summary: every project MUST use Podman rootless (or equivalent) for ALL containerized workloads — rootful Docker/sudo FORBIDDEN unless §11.4.112 platform constraint documented; `vasic-digital/containers` Submodule

(§11.4.76) sole orchestration layer via `pkg/boot` / `pkg/compose` / `pkg/health` ; extend upstream per §11.4.74 for missing capabilities; integration tests boot infra on-demand. Honest boundary: does NOT replace §11.4.10 credentials nor §12.3 hygiene. Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-161-PROPAGATION` + recommended gate `CM-ROOTLESS-CONTAINER-RUNTIME` + paired §1.1 mutation.

§11.4.162 — OpenDesign UI design system mandate (User mandate, 2026-06-21). Compact summary: every project with user-facing interfaces MUST use OpenDesign (<https://github.com/nexu-io/open-design>) as mandatory UI design system — NOT ad-hoc CSS; install as dependency, use token system for all color (light+dark from brand assets), typography, spacing, component tokens; extend upstream (§11.4.74) for unsupported patterns; every UI component ships light+dark + token-diff regression tests via visual regression; no element overlap/font collision/label overlay. Honest boundary: does NOT replace §11.4.27 functional testing nor §11.4.48/.49 UI testing. Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-162-PROPAGATION` + recommended gate `CM-OPENDESIGN-UI-SYSTEM` + paired §1.1 mutation.

§11.4.163 — Universal Media Validation (User mandate). Every recorded artifact MUST pass a MEDIA VALIDATION pipeline — OCR for video/screenshots, transcription for audio, text parsing; compare against SPECIFY-phase patterns (§11.4.159(J)); self-validated analyzer (§11.4.107(10)); structured verdict with pinpoint on FAIL; post-recording and real-time triggers. Paired §1.1 mutation ensures golden-bad fixture FAILs. Classification: universal. Propagation gate `CM-COVENANT-114-163-PROPAGATION` + paired §1.1 mutation.

§11.4.164 — Universal Auto-Propagation Hook (User mandate). Every constitutional pull triggers `post_update_hook.sh` — detects changed files, registers skills/MCP/hooks/scripts, logs summary. Consumer MUST invoke after every pull. Classification: universal. Propagation gate `CM-COVENANT-114-164-PROPAGATION` + paired §1.1 mutation.

§11.4.165 — Universal Independent Verification Agent (User mandate). Every code/media/docs/config output MUST pass INDEPENDENT verifier (§11.4.70/.20), iterate to GO per §11.4.134. CODE: review+build+mutations+runtime-signature.

MEDIA: pipeline+genuine+format. DOCS: exports+revision. CONFIG: syntax+schema+leaks. Self-validated by §1.1 mutation. Classification: universal. Propagation gate `CM-COVENANT-114-165-PROPAGATION` + paired §1.1 mutation.

§11.4.166 — REPEALED (operator decision, 2026-06-22). The Universal Semgrep static-analysis mandate is repealed — Semgrep is NO LONGER mandatory. Static analysis remains encouraged, not mandated. Scaffolding, `submodules/semgrep`, MCP/pre-commit/PATH wiring, `docs_chain` `semgrep_status` context, and the `CM-COVENANT-114-166-PROPAGATION` / `CM-SEMGREP-WIRED` gates removed. Anchor 11.4.166 retired. See constitution `Constitution.md` §11.4.166.

§11.4.167 — Big-work-item feature work-stream lifecycle (User mandate, 2026-06-23). Every BIG work item (new feature OR drastic/large fix) MUST be its own isolated feature work-stream — own sibling project copy, own `feature/<slug>` branch, own per-feature flashable builds + feature tags, SEPARATE from trunk until operator manually approves after testing, trunk merged INTO every stream regularly. (A) one big item = one sibling project copy (`<project>_<slug>`); small changes stay on trunk. (B) disk-feasible CoW/reflink clone MANDATORY (`cp -a --reflink=auto` on btrfs/XFS/ZFS/APFS, else `git worktree`); naive deep copies FORBIDDEN; per-stream `out/` + shared ccache; ~0-disk verified (§11.4.69) before relied on (§11.4.6). (C) own branch + greppable tags (§11.4.151 `<base>-feat-<slug>`), NEVER merged to trunk until §11.4.40 GREEN on clean target → §11.4.41 merge-first → §11.4.113 ff-only. (D) trunk merged INTO every live stream per trunk-tag/daily (merge, never rebase). (E) feature branch + tag cascades to EVERY touched submodule (§11.4.113). (F) single-builder build QUEUE (§11.4.58/.12.7/.12.8) + finite-device queue (§11.4.119); idle `out/` GC'd (§11.4.77). (G) every change crosses full gauntlet (§11.4.142/.125/.134/.145 + anti-bluff). (H) stream registry (§11.4.116/.147) + atomic `.partial` → rename + §11.4.84 resume + §9.2 retire-backup; bounded §12.12.6. Honest boundary (§11.4.6): isolation+disk-feasibility+greppable-identity+merge-discipline, NOT correctness (still §11.4.108/.40). Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-167-PROPAGATION` (literal `11.4.167`) + gates `CM-FEATURE-WORKSTREAM-COW-CLONE` / `-NO-MERGE-UNTIL-APPROVED` / `-TRUNK-SYNC-CADENCE` / `-SUBMODULE-CASCADE` + paired §1.1 mutations.

Canonical authority: constitution `Constitution.md` §11.4.167.

§11.4.168 — Exported-document independent content + textual + full-visual validation mandate (User mandate, 2026-06-23). Every generated/exported document (HTML/PDF/DOCX/any §11.4.65 export-scope format) MUST pass INDEPENDENT validation by a reviewer structurally separate from the generator (§11.4.70/.142, NEVER author self-check), on BOTH source AND every exported artifact, ALWAYS, across THREE layers: (1) CONTENT — export faithfully carries source intent+data, nothing dropped/truncated/garbled; (2) TEXTUAL — human-readable, NO raw markup/diagram-source (Mermaid `gantt` / `graph` / `flowchart` / `sequenceDiagram` etc.)/unrendered code-fence leaking as body text (the exact 2026-06-23 raw-gantt-in-PDF failure); (3) FULL VISUAL — diagrams render as IMAGES not source, layout intact, no overlap/cut-off/garble, verified by RENDERING the export (`pdftotext` catches raw source + `pdfimages` confirms rendered images + `pdftoppm` → OCR confirms human-readable content) with captured evidence §11.4.5/.69/.107. Anti-bluff: rubber-stamp ≠ validation; analyzer self-validated golden-good/golden-bad §11.4.107(10) (golden-bad seeded with raw `gantt` MUST FAIL); iterate to clean GO §11.4.134. Forensic FACT (2026-06-23): a “green” Mermaid fix shipped raw gantt source as readable text into user PDFs because validation grepped HTML for `class="mermaid"` and NEVER checked the exported PDF — operator caught it (§11.4.138). Honest boundary (§11.4.6): confirms the exported artifact is faithful/readable/visually-rendered — does NOT prove source content is correct (§11.4.142/.135) nor replace the §11.4.65 mtime freshness gate (this is the READABILITY/FIDELITY layer the mtime gate cannot see); FLAG_SECURE/un-rasterisable surfaces document the gap §11.4.112 + §11.4.117 proxy, never a faked PASS. Classification: universal (§11.4.17). Composes §11.4.65/.73/.107/.117/.135/.138/.142/.159/.163/.165/§1.1. Propagation gate `CM-COVENANT-114-168-PROPAGATION` (literal `11.4.168`) + recommended gate `CM-EXPORTED-DOC-VISUALLY-VALIDATED` + paired §1.1 mutation. **Canonical authority:** constitution `Constitution.md` §11.4.168.

§11.4.169 — Mandatory comprehensive test-type coverage with anti-bluff captured evidence (User mandate, 2026-06-25). Every project MUST cover a CLOSED enumerated test-type set, each to as-close-to-100% as the domain permits, every PASS citing rock-solid captured PHYSICAL evidence (§11.4.5/.69/.107), zero false results/zero bluff (§11.4.1/.6): (1) unit (only layer mocks allowed §11.4.27(A)); (2) integration (real System, infra via containers submodule §11.4.76); (3) e2e; (4) full-automation (autonomous, re-runnable, no manual step §11.4.25/.52/.98, deterministic §11.4.50); (5) Challenges (challenges

submodule banks §11.4.27(B)); (6) HelixQA (helix_qa submodule — written test banks/suites per app/service/platform + comprehensive fully-autonomous QA sessions §11.4.27); (7) DDoS/load-flood (clean refuse or graceful degrade); (8) security (authz/injection/secret-leak §11.4.10/crypto/CVE/default-deny + security submodule); (9) stress+chaos (load+contention+boundary+failure-injection, categorised recovery §11.4.85); (10) concurrency/atomicity (no lost updates, transactional atomicity, idempotent replay); (11) race-condition/deadlock (race detector + lock-order analysis; no blocking op in shared-lock region §1); (12) memory (leak census + peak-RSS ceiling + 24h/N-iter soak, no unbounded growth); (13) benchmarking/performance (p50/p95/p99 + throughput vs recorded baseline). The ONLY permitted absence is an honest §11.4.3 SKIP-with-reason, never a silent gap; four-layer §11.4.4(b) enforcement + coverage ledger (§11.4.25/.52); missing required type or OPERATOR_ATTENDED_ONLY = release blocker. Strict enumerated expansion of §11.4.27. Composes §11.4.4/.5/.6/.25/.27/.50/.52/.69/.76/.85/.98/.107/.123/.135/.146/§1.1. Classification: universal (§11.4.17). Propagation gate CM-COVENANT-114-169-PROPAGATION (literal 11.4.169) + recommended gate CM-MANDATORY-TEST-TYPES-COVERED + paired §1.1 mutation. **Canonical authority:** constitution Constitution.md §11.4.169.

§11.4.170 — Device-independent host-side rendered-UI visual-proof mandate (User mandate, 2026-06-25).

Compact summary: every change to ANY user-facing UI surface (screen/component/view/widget/layout/theme on any platform — Android-Compose, iOS-SwiftUI/UIKit, web, desktop, TUI) MUST be proven by DEVICE-INDEPENDENT host-side RENDERED PIXELS before it may be claimed correct — the real component rendered to a PNG ON THE HOST (NO device, NO emulator, NO running app) via a host-render harness (Compose → Roborazzi/Paparazzi on the JVM; web → Playwright/Storybook snapshot; SwiftUI → snapshot-testing; equivalent per stack), for EVERY relevant screen×state×{light,dark} theme, with the PNG validated by BOTH (i) golden image-diff (per-pixel/perceptual PASS/FAIL) AND (ii) an OCR/vision oracle reading rendered text+labels+control bounds to assert legibility + NO overlap / NO label-over-label / NO clipping / NO off-screen control / NO collapsed-or-giant-unbounded widget (the §11.4.162 layout invariants PROVEN on real pixels). Forensic FACT (2026-06-25): UI work was “validated” by VALUE-EQUALITY unit tests (a hex colour string / an sp/dp size integer / a design-token field == a constant) that NEVER RENDERED A PIXEL — GREEN while the operator opened a broken unbounded giant-button screen. A value/token-equality /

property-assertion UI test is FORBIDDEN as the PROOF a UI is correct (it may SUPPLEMENT, NEVER SUBSTITUTE the rendered-pixel proof); a “green” UI suite with no rendered-pixel artefact for the changed surface is a §11.4 PASS-bluff at the visual layer. Self-validated golden-good/golden-bad analyzer §11.4.107(10), thresholds calibrated on the project’s own fixtures §11.4.6. “device offline” is NEVER a valid skip — host-render IS the device-independent path; honest §11.4.3 SKIP only where the platform genuinely has no host-render harness (tracked migration item, never a value-only fallback). Honest boundary §11.4.6: host-render proves the COMPONENT renders correctly device-independently per commit — it does NOT prove the running app on real hardware (that remains §11.4.153/.158/.159/.160’s live-device recording + read-the-screen layer, which §11.4.170 COMPLEMENTS not replaces — both mandatory), does NOT replace §11.4.162 OpenDesign token/theme governance (§11.4.170 PROVES the rendered RESULT of those tokens), DISTINCT from §11.4.168 exported-document visual validation. Classification: universal (§11.4.17). Composes §11.4.3/.5/.27/.40/.69/.107/.108/.117/.153/.158/.159/.160/.162/.163/.168/.169/§1.1. Propagation gate CM-COVENANT-114-170-PROPAGATION (literal 11.4.170) + recommended gate CM-HOST-RENDERED-UI-VISUAL-PROOF + paired §1.1 mutation. **Canonical authority:** constitution submodule Constitution.md §11.4.170. Non-compliance is a release blocker. No escape hatch — no --value-assertion-proves-ui , --skip-rendered-pixel-proof , --token-equality-suffices , --device-offline-skip-visual , --single-theme-only , --unvalidated-image-diff-OK , --no-ocr-layout-oracle flag.

Companion documents

§11.4.171 — Mandatory comprehensive human-readable workable-item descriptions (User mandate, 2026-06-29). Every workable item (ATM-NNN ticket) in the project’s workable-items database, Issues.md, Fixed.md, Issues_Summary.md, Fixed_Summary.md, and ALL derived documents MUST carry a comprehensive human-readable description that explains the item in plain, non-technical language suitable for non-developers (project managers, stakeholders, team leads, business partners). The description MUST be 5-7 sentences minimum, written so that ANY person — regardless of technical background — can understand: (a) WHAT the item is (the feature, fix, or task in simple terms); (b) WHY it matters (the user/business value); (c) HOW it works (the

mechanism, without code references); (d) WHO benefits (the end user or stakeholder); (e) WHAT the expected outcome is (the measurable result). The description is MANDATORY — an item without a human-readable description is a §11.4 violation at the documentation layer. Descriptions MUST NOT use jargon, acronyms without expansion, code references, or technical implementation details as the primary explanation. Technical details MAY appear as supplementary context AFTER the plain-language explanation. The `description` column in the SQLite database (`docs/workable_items.db`) is the SINGLE SOURCE OF TRUTH for item descriptions; all derived documents (Issues.md, Fixed.md, summaries, exports) MUST carry the SAME description text. Pre-build gate `CM-HUMAN-READABLE-DESCRIPTION` (when implemented) walks every item in the DB and asserts the `description` field exists, is non-empty, and is ≥ 50 characters. Paired §1.1 meta-test mutation strips a description → gate FAILs. Classification: universal (§11.4.17). Composes §11.4.15 (item-status tracking), §11.4.16 (item-type tracking), §11.4.19 (column alignment), §11.4.91 (summary clarity), §11.4.93 (SQLite SSoT), §11.4.148 (item integrity), §11.4.65 (universal export). Propagation gate `CM-COVENANT-114-171-PROPAGATION` (literal `11.4.171`) + paired §1.1 mutation.

§11.4.172 — Mandatory production-readiness planning with realistic timeline projection (User mandate, 2026-06-29). Every project under this Constitution MUST maintain a living production-readiness planning document that: (a) provides REALISTIC timeline projections for all product milestones (MVP, beta, production-ready) based on actual development velocity data (items completed per week, reopening rate, blocking dependencies); (b) identifies ALL danger zones, weaknesses, and risk areas with mitigation strategies; (c) accounts for HW delivery timelines, licensing delays, and external dependency risks; (d) defines the critical path and identifies which items can slip without affecting MVP; (e) includes workstation/build-capacity analysis and optimization recommendations; (f) is updated at least monthly or whenever the item count changes by $\geq 10\%$; (g) is cross-referenced against the workable-items database for accuracy. The planning document MUST live at `docs/research/production_planning_<YYYYMMDD>/ANALYSIS.md` with HTML+PDF exports per §11.4.65. Honest boundary (§11.4.6): projections are estimates based on measured velocity — they guarantee the methodology is sound, NOT that specific dates will be met. Classification: universal (§11.4.17). Composes §11.4.6 (no-guessing — projections based on data, not

wishful thinking), §11.4.40 (full retest before tag), §11.4.42 (iteration discipline), §11.4.93 (SQLite SSoT), §11.4.108 (four-layer verification). Propagation gate `CM-COVENANT-114-172-PROPAGATION` (literal `11.4.172`) + paired §1.1 mutation.

§11.4.173 — Containerized + distributed build mandate (User mandate, 2026-06-29). EVERY build of EVERY component (source compile, artifact/package/installer/container-image production, codegen, asset render — for ANY language/platform: Go, Android/Gradle, desktop, web, native, firmware) MUST run INSIDE a specialized build container provisioned via the `digital.vasic.containers` submodule (§11.4.76) — NEVER on the bare host. The build containers MUST be DISTRIBUTED to the designated remote build host(s) (e.g. `thinker.local`) via the SAME containers-submodule distribution mechanism the infra uses (§11.4.76 Distributor / remote compose over SSH, §11.4.161 rootless), so the build EXECUTES on the remote build host (offloading the developer/main host); once the build completes the produced artifacts MUST be brought BACK to the originating main host (scp/rsync/volume copy) for use/flashing/distribution. Building outside a container, or on the bare host, is FORBIDDEN — a release blocker (the “works on my machine” / unreproducible-build class §11.4.76 exists to prevent). The build-host target + build-container definitions are config-injected (§11.4.28, never hardcoded in the submodule); a missing build-container capability is added by EXTENDING the containers submodule upstream (§11.4.74), never by an ad-hoc host build. Honest boundary (§11.4.6): the containerized+distributed build guarantees reproducibility + host-isolation + capacity offload — it does NOT replace the §11.4.40 full-suite retest, §11.4.108 four-layer artifact → runtime verification, or §11.4.38 installable-asset evidence (those still run against the brought-back artifact). Classification: universal (§11.4.17). Composes §11.4.76 (containers submodule — sole orchestration layer), §11.4.161 (rootless runtime), §11.4.74 (extend-don’t-reimplement), §11.4.28 (config injection / decoupling), §11.4.24 (build-resource stats), §11.4.82 (iteration speedup — persistent caches in the container), §11.4.121 (no-commit-while-build-writes-artifacts), §11.4.38 (installable-asset evidence on the brought-back artifact), §11.4.108 (artifact → runtime verification), §12.6 (host memory ceiling — offloading the build preserves the main host). Propagation gate `CM-COVENANT-114-173-PROPAGATION` (literal `11.4.173`) + recommended gate `CM-CONTAINERIZED-DISTRIBUTED-BUILD` (every build runs via the containers submodule on the remote build host; a bare-host build is detected + FAILs) + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md`

§11.4.173. Non-compliance is a release blocker. No escape hatch — no `--build-on-host` , `--skip-container-build` , `--local-build-ok` , `--no-distributed-build` , `--bare-host-build` flag.

§11.4.174 — Shared-host process-ownership verification mandate (User mandate, 2026-06-29). On any shared host where other projects/users/agents may run concurrently, every inspection of or action on a process, build, daemon, port, lock, or system-state signal MUST first positively VERIFY the target is OURS before diagnosing or acting. Attribute a process to us ONLY by a strong discriminator — cwd inside this project's tree, argv naming our scripts/artifacts/repo path, a PID we launched + recorded (§11.4.147), or a lock/port path under our tree — NEVER a loose `pgrep java/gradle/node` name match (it matches every project on the host). NEVER kill/signal/restart a process not positively ours (other projects' gradle/java/podman/zig are off-limits — operator-workload-safety, §12/§11.4.133); the §12.8 daemon-kill remediation MUST be scoped to OUR daemons. When our work waits on host contention from another project's process, surface it (§11.4.66/§11.4.101) — block-don't-break. Every `pgrep/ps` filter MUST exclude self-matches + other-project matches. Forensic anchor (FACT 2026-06-29): a 200%-CPU java+gradlew on the shared host was a separate `catalogizer` project's gradle test (cwd=Projects/catalogizer), nearly mis-attributed to our build. Composes §11.4.6/§11.4.66/§11.4.101/§11.4.147/§12/§12.8/§11.4.133. Classification: universal (§11.4.17). Propagation gate `CM-COVENANT-114-174-PROPAGATION` (literal `11.4.174`) + recommended gate `CM-PROCESS-OWNERSHIP-VERIFIED` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md`

§11.4.174. Non-compliance is a release blocker. No escape hatch — no `--assume-process-ours` , `--loose-pgrep-ok` , `--kill-any-matching-process` , `--skip-ownership-check` flag.

§11.4.176 — Conflict-free multi-track work-division + exactly-once claim registry + capability-aware deadlock-proof device-lock (User mandate, 2026-07-02). Two mandatory multi-track arbitration layers, conflict-free by construction: (A) exactly-once work-item/logical-group claim registry (extends §11.4.147/§11.4.116; disjoint file-scope §11.4.58 L3 + worktree/CoW §11.4.167; no codebase loss §9.2/§11.4.113/§11.4.84/§11.4.41); (B) capability-aware deadlock-proof device-lock (§11.4.119 single-owner; all-or-nothing + non-blocking + TTL-reap breaks all 4 Coffman conditions; append-only JSONL + atomic snapshot §11.4.116); (C) universal/decoupled (§11.4.17/§11.4.28/§11.4.35) + evidence-based resource-tuning (§11.4.24/§11.4.6, never guessed). Classification: universal

(§11.4.17). Gates CM-WORK-DIVISION-EXCLUSIVE-CLAIM + CM-DEVICE-LOCK-DEADLOCK-FREE + CM-COVENANT-114-176-PROPAGATION + paired §1.1 mutation. Composes §11.4.6 §11.4.17 §11.4.24 §11.4.28 §11.4.35 §11.4.41 §11.4.58 §11.4.84 §11.4.85 §11.4.103 §11.4.113 §11.4.116 §11.4.119 §11.4.147 §11.4.167 §12.6 §9.2 §11.4.176.

§12.11 — Maximal dynamic resource utilization for containerized builds (User mandate, 2026-07-03). The containerized build path (its OWN cgroup + own MemoryMax → OOM-kills ITSELF, never user.slice) MUST use the maximal SAFE fraction of host capacity, computed DYNAMICALLY per host: auto-detect nproc/MemTotal/RLIMIT_NPROC (never hardcode, §11.4.6), reserve explicit host headroom (cores+RAM), scale -j against BOTH the memory model AND the process rlimit (privileged nproc-raise for full -j; EAGAIN-safe -j otherwise), never break/bottleneck/wedge. NO fixed low -j for the containerized path. The §12.6 60% ceiling + bare-metal -j cap REMAIN, SCOPED to user.slice-resident work — §12.11 does NOT weaken §12.6, it scopes it. Classification: universal (§11.4.17). Gate CM-MAXRES-DYNAMIC-BUILD + CM-COVENANT-12-11-PROPAGATION + paired §1.1 mutation. Composes §12.1 §12.2 §12.3 §12.6 §12.10 §11.4.6 §11.4.24 §11.4.133 §9.2 §11.4.35 §12.11.

§11.4.177 — Developer-tooling project-decoupling + invocation-directory operation (User mandate, 2026-07-04). Compact summary: no project-specific script / hook / alias / binary may be wired (copied / symlinked / PATH-exported / aliased) into a global or shared developer-tooling PATH; shared tooling MUST be project-agnostic and operate on the INVOCATION directory (cwd / explicit path / config arg), NEVER a hardcoded project path; a project-specific helper lives in + runs from its own repo, a genuinely-reusable helper is promoted to the shared toolkit ONLY after being stripped of every project-specific assumption (paths, device serials, package names, region endpoints) + taking its target from the invocation context; a project script reachable on the global/shared PATH is a decoupling violation of equal severity to importing project-specific code into a shared submodule (§11.4.28); forensic FACT 2026-07-04 — a project cwd-hook symlinked into the global Claude-Toolkit PATH re-coupled every alias to one hardcoded checkout; honest boundary §11.4.6 — decoupling guarantees project-agnostic, not correct (still needs the tool's own tests). Classification: universal (§11.4.17). Composes §11.4.28/.29/.35/.11/.6. Propagation gate CM-COVENANT-114-177-PROPAGATION (literal 11.4.177) + recommended gate CM-TOOLING-PROJECT-DECOUPLED + paired §1.1 mutation. **Canonical authority:** constitution submodule

Constitution.md §11.4.177. Non-compliance is a release blocker. No escape hatch — no `--global-symlink-ok` , `--hardcode-project-path` , `--wire-into-shared-path` flag.

§11.4.178 — Track-qualified identity for parallel work streams (session-name collision ban) (User mandate, 2026-07-04). Compact summary: parallel work streams sharing hardware / a git backing store / a tmux-session namespace / a lock namespace / any single-owner resource MUST be addressed by a TRACK-QUALIFIED identity (`<project>__<track>__<role>`), NEVER a bare directory basename; multiple checkouts/clones/worktrees sharing basename `atmosphere` collapse into one session/lock/log key + cross-wire tmux targets, lock files, device claims, log dirs (observed collision 2026-07-04); every registry row / session name / lock path / log dir / device claim for a ≥ 2 -concurrent-claimant resource MUST carry the track-qualified key, a bare-basename identity for such a resource is a violation; honest boundary §11.4.6 — a unique identity prevents cross-wiring, it does not itself arbitrate ownership (that is §11.4.119 + **§11.4.176** (the multi-track work-division + exactly-once-claim + device-lock arbitration anchor)). Classification: universal (§11.4.17). Composes §11.4.111/.119/.29/.58 + **§11.4.176** (the multi-track work-division + exactly-once-claim + device-lock arbitration anchor). Propagation gate `CM-COVENANT-114-178-PROPAGATION` (literal `11.4.178`) + recommended gate `CM-TRACK-QUALIFIED-IDENTITY` + paired §1.1 mutation. **Multi-project-per-track layout clarification (2026-07-10):** a `/mnt/track<N>` track parent is NOT single-project — MULTIPLE projects coexist on the same track parents, each in its OWN project-named subdirectory `/mnt/track<N>/<project_name>/` (e.g. `/mnt/track1/atmosphere/` AND `/mnt/track1/helix_ota/` side by side); the per-track working copy of project P on track N is ALWAYS `/mnt/track<N>/<project_name>/` (canonical lowercase snake_case dir-name per §11.4.29), never the bare track root and never a shared basename; this is WHY the identity carries the `<project>` component (`<project>__<track>__<role>`) — so two coexisting projects' track-N sessions / locks / logs / device-claims / registry rows never collide; a project adopting the multi-track engine (§11.4.187) MUST create AND register its own `/mnt/track<N>/<project_name>/` per track and MUST NOT assume it owns the track root or that a track hosts only one project.

Canonical authority: constitution submodule Constitution.md §11.4.178. Non-compliance is a release blocker. No escape hatch — no `--bare-basename-ok` , `--skip-track-qualifier` , `--shared-session-name` flag.

§11.4.179 — Corruption-isolated parallel git streams (own-.git independent clones, not shared-common-dir worktrees) (User mandate, 2026-07-04).

Compact summary: when corruption-isolation is a requirement, parallel git work streams MUST each be an isolated repo with its OWN `.git` (own object store, own index, own lock namespace), NOT `git worktree` checkouts sharing one common `.git`; a shared common-dir is a single point of failure — one stale `index.lock` / `.commit_all.lock` / `.push_all.lock` freezes EVERY stream at once (observed 9-h all-track freeze 2026-07-04) + one corrupted object store loses every stream; each stream's `git rev-parse --git-common-dir` MUST resolve INSIDE its own tree, a common-dir resolving to a shared parent is a violation; where CoW/reflink disk-feasibility is the constraint use the §11.4.167 reflink-clone-with-own-`.git` path (btrfs/XFS/ZFS/APFS reflink shares extents on disk while keeping a per-stream `.git`, so isolation + disk-feasibility both hold); honest boundary §11.4.6 — own-`.git` contains lock/corruption blast radius, it does NOT remove per-repo stale-lock reaping (§11.4.180) nor ff-only no-force integration (§11.4.113). Classification: universal (§11.4.17). Composes §11.4.167/.119/.58/§9.2 + **§11.4.176** (the multi-track work-division + exactly-once-claim + device-lock arbitration anchor). Propagation gate `CM-COVENANT-114-179-PROPAGATION` (literal `11.4.179`) + recommended gate `CM-GIT-STREAMS-OWN-COMMON-DIR` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.179. Non-compliance is a release blocker. No escape hatch — no `--shared-common-git`, `--worktree-when-isolation-required`, `--skip-git-isolation` flag.

§11.4.180 — Commit/push (single-writer) wrappers MUST auto-reap provably-stale locks before acquiring (User mandate, 2026-07-04).

Compact summary: every commit / push / single-writer wrapper MUST, before acquiring its own lock, auto-reap any PROVABLY-stale git lock — one whose recorded holder PID is DEAD (`kill -0` fails), OR (no PID recorded) older than a defined threshold AND with no live wrapper/git process holding that repo; the reap covers the wrapper's own lock plus the git-internal existence-based locks (`index.lock`, `HEAD.lock`, `refs/**/*.*.lock`) git does NOT auto-release on holder death; it MUST NEVER remove a lock whose holder is ALIVE (removing a live-held lock corrupts a concurrent writer — §9.2); each reap decision (REAPED / KEPT + reason) is logged as captured evidence (§11.4.6 — liveness PROVEN via `kill -0`, never assumed); a wrapper that blocks indefinitely on a dead-holder lock (observed 9-h freeze 2026-07-04) is a violation; honest boundary §11.4.6 — reaping clears dead-

holder deadlock, it does NOT substitute for the own- .git isolation of §11.4.179. Classification: universal (§11.4.17). Composes §11.4.84/.88/.6/§9.2/§11.4.116/§11.4.179. Propagation gate `CM-COVENANT-114-180-PROPAGATION` (literal `11.4.180`) + recommended gate `CM-WRAPPER-STALE-LOCK-AUTO-REAP` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.180. Non-compliance is a release blocker. No escape hatch — no `--skip-stale-lock-reap`, `--force-remove-live-lock`, `--block-on-dead-holder` flag.

§11.4.181 — Consistent controlled feature-branch naming: one feature/logic-group ⇒ exactly ONE canonical branch name across the main repo + all owned submodules (User mandate, 2026-07-05). Compact summary: one feature / logical group of workable items MUST map to EXACTLY ONE canonical branch name used IDENTICALLY on the main repo AND every touched owned submodule — NEVER two differently-named branches for one feature/group, NEVER per-repo naming drift; (1) single canonical name per group minted ONCE per the §11.4.167 D scheme (`feature/<slug> branch + <prefix>-<base>-feat-<slug> tag, <slug> lowercase snake/kebab §11.4.29`), used identically across the main repo + every branched submodule (an untouched submodule stays on its base branch); (2) single source of truth — the canonical name is recorded ONCE in the §11.4.176 exactly-once claim registry / the §11.4.93 workable-items `logic_group → branch` mapping and LOOKED UP, never re-invented (creating a branch whose name ≠ the group's registered canonical name, or minting a 2nd name for a group that already has one, is a §11.4.6 no-guessing violation); (3) mechanical enforcement (§11.4.75) — recommended gate `CM-BRANCH-NAME-CONSISTENCY` FAILs on an unregistered feature-branch name / two divergent names for one group / a submodule branch name diverging from the group's main-repo canonical name, paired §1.1 mutation (register one group under two divergent names → gate FAILs); (4) risk-free reconciliation (§9.2/§11.4.113/§11.4.84) — a mis-named/duplicate branch is reconciled by MERGING its commits onto the canonical branch (union, no commit lost, no `-s ours /reset`), ff-only push to all upstreams (NEVER force-push), retiring the mis-named branch only after the merge is confirmed on all mirrors. Classification: universal (§11.4.17). Composes §11.4.167/.176/.178/.179/.28/.29/.93/.113/.84/.6/.75/§9.2. Propagation gate `CM-COVENANT-114-181-PROPAGATION` (literal `11.4.181`) + recommended gate `CM-BRANCH-NAME-CONSISTENCY` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.181. Non-compliance is a

release blocker. No escape hatch — no `--allow-divergent-branch-names` , `--skip-branch-registry` , `--rename-branch-freely` , `--per-repo-branch-name-OK` , `--two-names-one-feature-OK` flag.

§11.4.182 — Track+branch work-stream identity label (T<N>/<branch>) on every agent/subagent label + every operator-facing work-stream reference (User mandate, 2026-07-05). Compact summary: EVERY agent / subagent / work-stream label AND every operator-facing reference to a work stream MUST START with a (T<N>/<branch> - <alias>) prefix identifying the track + git branch + the claude-toolkit alias in use (e.g. (T1/main - claude1) ATM-312 MPV drm_prime investigation , (T2/feature/mistiq-vader - deepseek) SPK-XXX ...) so parallel-track work (/mnt/track1..4 , each on its own branch, each driven by a distinct toolkit alias) is never ambiguous; (1) universal label form on every agent-dispatch description, status report, resume/handoff doc, progress line; (2) DETERMINISTIC derivation never guessed (§11.4.6) — <N> from the checkout path /mnt/track<N>/... (honest ? off-track, never a guessed number), <branch> from git rev-parse --abbrev-ref HEAD , <alias> from CLAUDE_CONFIG_DIR basename .claude-<alias> → <alias> (honest ? when unset/non-matching, never a guessed name) (reference labeler scripts/multitrack/track_branch_label.sh); (3) mechanical enforcement (§11.4.75 / §11.4.109-class) — a PreToolUse guard hook (constitution/scripts/hooks/guard-track-branch-label.sh , inherited BY REFERENCE per §11.4.177, never copied) BLOCKS any Agent/Task/TaskCreate dispatch whose description (or subagent label) does not START with ^\([0-9]+\|^)+\ , non-agent tools untouched; pre-build gate CM-TRACK-BRANCH-LABEL (labeler+hook exist/exec/parseable, hook wired in settings, convention doc exists) + propagation gate CM-COVENANT-114-182-PROPAGATION (literal 11.4.182); (4) honest boundary (§11.4.6) — a ? in the track OR the alias field is a valid honest label, never a bluff, and the enforced numeric-track format surfaces (blocks) an off-track dispatch rather than mislabeling it. Classification: universal (§11.4.17). Composes §11.4.178 (track-qualified identity — the label-prefix instantiation applied to every agent/subagent label + operator-facing reference) / §11.4.29 / §11.4.75 / §11.4.109 / §11.4.177 / §11.4.6 / §11.4.58 / §11.4.103 / §11.4.119 / §11.4.116 / §11.4.147. Propagation gate CM-COVENANT-114-182-PROPAGATION (literal 11.4.182) + recommended gate CM-TRACK-BRANCH-LABEL + paired §1.1 mutation. **Canonical authority:** constitution submodule Constitution.md §11.4.182. Non-compliance is a

release blocker. No escape hatch — no `--skip-track-label` , `--unlabeled-agent-OK` , `--bare-work-stream-reference` , `--no-track-branch-prefix` , `--label-from-memory` flag.

§11.4.183 — Maximal multi-agent utilization per work-stream + full-constitution-application + zero-bluff mandate (User mandate, 2026-07-07).

Compact summary: every track / work-stream MUST maximize multi-agent working approaches — subagent-driven development (§11.4.20/§11.4.70), independent review agents (§11.4.125/§11.4.142/§11.4.165/§11.4.134), parallel background streams with auto-backfill (§11.4.58/§11.4.103), and every other applicable agent form — so every track produces the MOST completed work at the HIGHEST quality in the LEAST time; every track MUST apply the ENTIRE constitution (every rule, mandatory constraint, guide, and derived technology — NOTHING ignored, skipped, avoided, or deemed less-relevant); there MUST NEVER be false / faulty / untested / unverified results or bluff of any kind anywhere (§11.4 anti-bluff covenant, captured evidence §11.4.5/§11.4.69/§11.4.107); honest boundary §11.4.6 — “maximize agents” is bounded by §12.6 60% memory + §11.4.58 parallel-agent cap + genuine parallelizability + §11.4.119 single-resource-owner, spawning contending/redundant agents is NOT maximization, the bar is maximum USEFUL parallel throughput not agent count. Classification: universal (§11.4.17). Composes §11.4.20/.58/.70/.92/.94/.97/.103/.119/.125/.126/.134/.142/.165/§12.6. Propagation gate `CM-COVENANT-114-183-PROPAGATION` (literal `11.4.183`) + recommended gate `CM-MAX-AGENTS-PER-TRACK` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.183. Non-compliance is a release blocker. No escape hatch — no `--single-agent-OK` , `--skip-review-agents` , `--partial-constitution-application` , `--serialise-parallelisable-work` flag.

§11.4.184 — Mandatory SonarQube static-analysis CLI + local tooling installed and PATH-discoverable (User mandate, 2026-07-06).

Compact summary: every project/host that runs static analysis MUST have the SonarQube scanner CLI (`sonar-scanner`) installed AND durably discoverable on the system PATH via the shell rc (`.bashrc` / `.zshrc`) — exactly like the Firebase/GitHub/GitLab CLIs — PLUS the shared `constitution/scripts/sonarqube/` tooling (`sonarqube_install_check.sh` / `sonarqube_lib.sh` / `sonarqube_run_scan.sh` / `sonarqube_container.sh` + `compose/`) consumed BY REFERENCE (§11.4.28 decoupled + §11.4.177 project-agnostic + §11.4.80-style inherited, NEVER copied per project); (1) installed + PATH-durable — the

scanner resolves in every fresh shell via a durable rc export, not an ephemeral session-only PATH; (2) shared tooling present + GREEN — `sonarqube_install_check.sh` exits 0 (scanner + rootless podman §11.4.161 + podman-compose + curl + ES-ready `vm.max_map_count`) before any scan-dependent work; (3) anti-bluff §11.4.6 — “installed/configured” PROVEN by `command -v sonar-scanner` resolving AND install-check exit 0, NEVER assumed, a claimed-installed-but-absent CLI is a §11.4 tooling-layer bluff; (4) rootless §11.4.161 — the local SonarQube server runs via rootless podman, never rootful docker/sudo; honest boundary §11.4.6 — names the SonarQube SCANNER CLI + shared local-server tooling, and per §11.4.166 (former universal Semgrep static-analysis mandate REPEALED) SonarQube is a DISTINCT operator-mandated tool not a Semgrep re-instatement, mandating TOOLING availability not a scan on every build. Classification: universal (§11.4.17). Composes §11.4.28/.36/.74/.75/.109/.161/.166/.177/.6. Propagation gate `CM-COVENANT-114-184-PROPAGATION` (literal `11.4.184`) + recommended gate `CM-SONARQUBE-CLI-INSTALLED` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.184. Non-compliance is a release blocker. No escape hatch — no `--skip-sonarqube-install` , `--sonar-cli-optional` , `--ad-hoc-sonar-path` , `--rootful-sonar-OK` flag.

§12.12 — Process/thread-limit (RLIMIT_NPROC) awareness for parallel subagent/multi-process work (User mandate, 2026-07-07). Compact summary: heavy parallel subagent/multi-process work is bounded by the OS per-user process/thread limit `ulimit -u` (`RLIMIT_NPROC` , Linux counts ALL threads+processes of the real UID, not per-process); the ambient tooling fleet (MCP servers, tokio/V8/libuv worker runtimes, `mongod` / `chrome` / `prisma` -class MCP servers, local LLM servers) consumes a large FIXED fraction before project work starts — FORENSIC FACT (2026-07-07): a host at `ulimit -u 4096` sat at ~3900-4005 threads from the ambient fleet alone (<150 headroom); a 4th Go-heavy subagent (`GOMAXPROCS=nproc` spawns dozens of threads) hit `EAGAIN` / `failed to create new OS thread (errno=11)` and crashed tooling; stopping 3 non-essential MCP servers freed ~1345 threads (3827 → 2482). Mandate: BEFORE scaling parallel subagents/processes, check thread headroom (`ulimit -u` `soft` + `ulimit -Hu` `hard` + `live ps -L --no-headers -u "$USER" | wc -l`) and treat exhaustion as a §12 host-safety event that YIELDS UNCONDITIONALLY — the ORTHOGONAL second exhaustion axis alongside §12.6’s memory ceiling (abundant RAM ≠ abundant thread budget); signals `EAGAIN` / `fork: retry` /

failed to create new OS thread / cannot allocate memory on fork/clone (distinct from OOM, §11.4.6 no-guessing — never conflate); low headroom ⇒ SERIALIZE, avoid `GOMAXPROCS=nproc` -class spikes, never dispatch blind (a subagent hitting `EAGAIN` mid- git push is a §9.2 data-safety risk). Three raising techniques: (a) no-restart `prlimit --pid <PID> --nproc=<soft>:<hard>` (root, e.g. `su -c`) — new children of a running process inherit the raise; (b) persistent `/etc/security/limits.conf` (`<user> hard/soft nproc <N>`) or `systemd` `LimitNPROC=<N>`, next login; (c) self-raise `ulimit -u <N>` up to the existing hard ceiling within the current session — `su -c "ulimit -u N"` sets the limit ONLY inside that transient subshell, NOT the already-running session. Freeing threads (stopping non-essential MCP servers) REQUIRES operator authorization (§11.4.122, no autonomous tooling removal); CAUTION `pkill -f` / `pgrep -f` match the killer's OWN command line — always exclude the caller's `$$` / `$PPID` (and conductor PID) from any cmdline-pattern kill list to avoid self-kill.

Classification: universal (§11.4.17). Composes §12/.6/.7/§11.4.58/.101/.103/.122/ §9.2. Propagation gate `CM-COVENANT-12-12-PROPAGATION` (literal `12.12`) + recommended gate `CM-NPROC-HEADROOM-CHECK` + paired §1.1 mutation.

Canonical authority: constitution submodule `Constitution.md` §12.12. Non-compliance is a release blocker. No escape hatch — no `--ignore-thread-limit`, `--dispatch-blind`, `--skip-headroom-check`, `--autonomous-tooling-kill`, `--assume-ulimit-scope` flag.

§11.4.185 — Manual QA-team testing as the mandatory FINAL confirmation of done (User mandate, 2026-07-07). Compact summary: no set of features / scope of work / release may be considered fully completed until it has received MANUAL TESTING confirmation by the project's QA team as the FINAL confirmation step — everywhere and always, never forgotten; every automated gate (§11.4 anti-bluff covenant, §11.4.5/§11.4.69 captured evidence, §11.4.52 autonomous validation, §11.4.40 full-suite retest, §11.4.108 four-layer verification, §11.4.132 risk-ordered validation) is NECESSARY but NOT SUFFICIENT — manual QA is the FINAL human sufficiency gate; ADDS to (never weakens) §11.4.52's mandatory autonomous-validation-path requirement — a feature needs BOTH the autonomous path AND the manual-QA final confirmation, never one substituting the other; pipeline: autonomous-GREEN → build → flash/deploy → autonomous on-device validation → hand off to QA team → ONLY after QA-team manual confirmation is the scope fully-completed / tag-eligible; the §11.4.126 autonomous-loop release-tag terminal condition now REQUIRES this confirmation; honest boundary §11.4.6 —

the agent MUST hand off + WAIT, never self-certify/simulate/infer the manual step, while continuing every other non-blocked item in parallel per §11.4.87/.94/.97/.101 (the wait parks one scope, not the loop). Classification: universal (§11.4.17). Composes §11.4/.5/.6/.40/.52/.69/.108/.126/.129/.130/.132. Propagation gate `CM-COVENANT-114-185-PROPAGATION` (literal `11.4.185`) + recommended gate `CM-MANUAL-QA-FINAL-CONFIRMATION` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.185. Non-compliance is a release blocker. No escape hatch — no `--skip-manual-qa`, `--automation-suffices`, `--self-certify-manual-step`, `--assume-qa-confirmed`, `--autonomous-only-completion` flag.

§11.4.186 — Anti-divergence enforcement: cross-document consistency is a mandatory-before-export/commit gate, never an after-the-fact audit (research-derived, 2026-07-08). Compact summary: any project maintaining more than one representation of the same tracked data (a delivery plan + its MVP brief + a workable-items DB + their summaries + `.md` / `.html` / `.pdf` / `.docx` exports) MUST enforce cross-document CONSISTENCY as a deterministic gate that runs BEFORE any export render, any doc/DB sync `verify`, and any doc-set commit — NEVER as an after-the-fact human audit that runs only when an operator happens to ask (the forensic FACT: the same NanoKVM/phone-as-screen feature appeared in three plan clusters — one shared ticket SPK-596 across SIX rows split over TWO releases, and a carve-out due to ship 2026-08-31 whose own prerequisite does not START until 2026-10-05 — and NO gate caught it; only an operator's manual read of `Plan_v9.0.xlsx` surfaced it, after the divergent bytes had already exported). (1) **One no-divergence verdict, three seams** — a single deterministic PASS/FAIL/SKIP verdict (built by composing a cross-document integrity validator with the workable-items-DB internal `validate`) gates all three write-seams (export / sync-verify / commit); a FAIL at any seam refuses that seam, naming the exact offending records (§11.4.6 actionable). (2) **Five decidable check families** — DEDUP (no two records describe one feature with divergent timeline/status/release, keyed on ticket OR normalised `(subject, scope)`, NOT bare subject substring), TIMELINE (start ≤ deadline; no deadline-before-dependency; no dependency cycle; GATED exempt-but-marked), CROSS-DOC (the same item across every representation carries identical timeline/status/type against a NAMED authoritative source), INTEGRITY (no orphan refs; Status ↔ Type §11.4.33; location ↔ status), STRUCTURAL (required columns; ID uniqueness+monotonicity §11.4.54). (3) **Drift-proof trigger** — a §11.4.86 sha256-of-sorted-members fingerprint of the

authoritative inputs re-arms the gate on ANY input change, so a stale verify cannot pass on changed data. (4) **Self-validated analyzer (§11.4.107(10))** — the validator ships golden-good (MUST PASS) + one golden-bad per family (MUST FAIL, pinpointing the offender) + a negative-control (distinct same-subject tasks that MUST PASS — the false-positive guard); a validator that passes its golden-bad, or fails golden-good/the negative-control, is itself a §11.4 bluff and a release blocker. (5) **Reusable + decoupled (§11.4.28/.31)** — packaged as a depth-1 reusable engine under the constitution submodule per the §11.4.28(C) carve-out with a `helix-deps.yaml` ; project-specific doc-sets + authoritative-source bindings live in a consumer-owned checkset (data, never engine code). (6) **Supersedes ad-hoc audits** — the per-cycle `*_INTEGRITY_FINDINGS.md` / `RECONCILIATION_*.md` after-the-fact audit docs are retired in favour of the gate; honest boundary §11.4.6 — the gate proves internal CONSISTENCY (no record diverges, no timeline is incoherent, every ref resolves), NOT plan CORRECTNESS (achievability/date-rightness stay operator/management decisions the gate SURFACES, never MAKES). Classification: universal (§11.4.17). Composes §11.4.6/.12/.33/.44/.50/.53/.54/.60/.65/.73/.75/.85/.86/.91/.93/.95/.106/.107/.108/.110/.148/.176. Propagation gate `CM-COVENANT-114-186-PROPAGATION` (literal `11.4.186`) + recommended gate `CM-DOC-INTEGRITY-VALIDATION` (5 invariants: validator present+executable; consumer checkset present; the no-divergence gate wired into each of the three seams; fingerprint sidecar fresh; `selfcheck` golden-good/golden-bad/negative-control wired into meta-test) + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.186. Non-compliance is a release blocker. No escape hatch — no `--skip-divergence-gate` , `--audit-after-export-OK` , `--cross-doc-check-not-applicable` , `--unvalidated-analyzer-OK` , `--divergent-commit-OK` flag.

§11.4.187 — Automatic multi-track ruler orchestration (User mandate, 2026-07-04; landed 2026-07-09). Compact summary: every project incorporating this constitution MUST ship its multi-track parallel-development orchestration as a UNIVERSAL, AUTOMATIC, OUT-OF-THE-BOX, INHERITED capability — not a per-project reimplementation — via a single conductor session that programmatically SPAWNS, DRIVES, RESUMES, MONITORS, and CONTROLS per-track headless workers. (1) **Headless spawn primitive** — workers launch via native headless `claude -p --output-format stream-json --verbose` ; the worker's `session_id` is captured from the FIRST `init` event (never invented);

success/failure is read from the terminal event's `result.is_error` field, NEVER the process exit code (a CLI can exit non-zero on a successful `result`, or zero on a failed one — exit-code-as-success-signal is a §11.4.6 guess); continuation of an existing worker uses `--resume <session_id>` invoked from the SAME cwd + env the spawn used (the storage/resume-collision gotcha). (2) **Per-subscription auth + rate-limit fallback** — every spawn `unset` s any ambient `ANTHROPIC_API_KEY` (the precedence trap that silently routes a “subscription” worker onto pay-per-token API billing) and sets the selected alias's `CLAUDE_CODE_OAUTH_TOKEN` / `CLAUDE_CONFIG_DIR`; on a captured rate-limit signature (`isApiErrorMessage + apiErrorStatus:429`, quota-exhausted vs transient distinguished by evidence, never guessed) the ruler REBINDS the affected track onto the next healthy alias (cross-track reuse permitted), falls back to a provider alias where configured, and bounded-parks (§11.4.101) rather than hanging when every alias is cooled. (3) **Crash-resilient ruler self-supervisor** — the ruler's OWN state (per-track alias/session_id/worktree/last-verdict) is persisted durably (write-temp-then-rename, NEVER tmpfs) so an external watchdog detects a dead ruler PID and re-hydrates every track's binding from the durable state + the §11.4.116 event stream — no work is silently lost to a ruler crash (§11.4.147 at the ruler layer). (4) **Idempotent bootstrap + auto-install + conductor-stays-home** — one idempotent bootstrap command wires the whole system out of the box (cwd-hook install, config validation, reconciliation), auto-invoked on every constitution pull via the §11.4.164 `post_update_hook.sh` skill-register seam; the `conductor`: -designated alias/session is NEVER bound into a track worktree (the resolver returns nothing for it) so the conductor always runs from the main checkout. (5) **Project-agnostic engine, consumer-owned data** — the entire mechanism (config loader, alias → worktree resolver, cwd-hook, bind/fallback/cooldown orchestrator, session launcher, rate-limit monitor, supervisor) lives at `constitution/scripts/multitrack/`, inherited BY REFERENCE (§11.4.28(B)/§11.4.177), carrying NO project literal; every consuming project supplies its own per-host `config/multitrack/<hostname>.yaml` (tracks/mounts/serials/aliases/conductor) as DATA, never as an engine edit. Anti-bluff (§11.4/§11.4.108): every piece carries four-layer coverage (pre-build gate + on-host test + paired §1.1 mutation + captured evidence per §11.4.69), and each fix declares a runtime signature verified on a clean host — never a claim the mechanism “should work.” Honest boundary (§11.4.6): this anchor mandates the MECHANISM; genuine per-subscription quota-isolation proof requires live tokens exercising real rate-limit

boundaries (an operator-scheduled live acceptance window) — a mock-driven GREEN proves the spawn/resume/fallback CONTRACT, never live-quota behaviour, and the two are never conflated or substituted for each other. Classification: universal (§11.4.17). Composes §11.4.20 (subagent-driven-by-default) / §11.4.28 (submodule decoupling + dependency layout) / §11.4.58 (parallel-development PWU pipeline) / §11.4.70 (subagent-driven execution default) / §11.4.101 (autonomous-decision-over-blocking — the bounded-park-on-all-cooled behaviour) / §11.4.103 (continuous parallel-stream routine) / §11.4.111 (resolve-by-stable-name — config by serial, never by index) / §11.4.116 (real-time conductor ↔ framework sync channel) / §11.4.119 (single-resource-owner partitioning) / §11.4.126 (default autonomous-loop mode) / §11.4.128 (always-on device-recording — the analogous host-safety discipline) / §11.4.131 (standing session-resumption file) / §11.4.147 (crashed-agent respawn-until-complete — the ruler-self-supervisor is its ruler-layer instantiation) / §11.4.164 (constitution auto-propagation hook seam) / §11.4.167 (feature work-stream lifecycle) / §11.4.176 (multi-track work-division + exactly-once claim + device-lock) / §12.6/§12.8 (host-safety memory + concurrency ceilings — the spawn/launch pool is budget-gated). Propagation gate `CM-COVENANT-114-187-PROPAGATION` (literal `11.4.187`) + recommended gates `CM-MULTITRACK-ENGINE-IN-CONSTITUTION` (the generic ruler scripts present under `constitution/scripts/multitrack/`, carrying no project literal, `bash -n + sh -n` clean per §11.4.67) / `CM-MULTITRACK-AUTOINSTALL-WIRED` / `CM-MULTITRACK-BOOTSTRAP-IDEMPOTENT` / `CM-MULTITRACK-BIND-ON-START` / `CM-MULTITRACK-FALLBACK-MONITOR` / `CM-MULTITRACK-HEADLESS-SPAWN` / `CM-MULTITRACK-RULER-SUPERVISOR` + paired §1.1 mutation (strip the literal → propagation gate FAILs; re-introduce a project literal into the engine, or downgrade a rate-limit rebind to exit-code-based success detection → the corresponding mechanism gate FAILs). **Canonical authority:** constitution submodule `Constitution.md` §11.4.187. Non-compliance is a release blocker. No escape hatch — no `--skip-multitrack`, `--manual-multitrack-only`, `--no-auto-install`, `--per-project-reimpl-OK`, `--exit-code-is-success-signal` flag.

§11.4.188 — Regular main → feature merge cadence: every track / feature-branch MUST FREQUENTLY merge canonical `main` into itself DURING the work, never only at the end (User mandate, 2026-07-09). Compact summary: every long-lived feature branch AND every parallel-development track (§11.4.58/ §11.4.103 parallel streams; §11.4.167 feature work-streams; §11.4.176/§11.4.178/

§11.4.179 multi-track) MUST regularly `git merge origin/main` (the canonical trunk/master) INTO its own branch THROUGHOUT the work — NOT only at the end — so no branch ever drifts far from main and the eventual back-merge stays a small, low-conflict, low-risk operation. GENERALISES §11.4.167(D) (trunk-merged-INTO-every-BIG-feature-STREAM, at minimum after every trunk tag / daily, `git merge` never rebase, stale-if-not-synced) from the big-work-item feature-stream lifecycle to EVERY feature branch on EVERY track — a small/medium feature branch, or any `git worktree` /CoW track that never earned a full §11.4.167 stream, is STILL bound to this cadence (the §11.4.167(D) discipline promoted to the general standalone rule; §11.4.167(D) remains its big-item specialisation). (1) **WHAT** — fetch the latest canonical `main / master` and merge it INTO the working branch, keeping the branch continuously close to trunk; a branch that has NOT integrated trunk within the cadence is flagged STALE (the exact long-lived-branch conflict-storm + integration-risk anti-pattern this forbids). (2) **CADENCE** — merge trunk in FREQUENTLY: after EVERY trunk tag, at minimum DAILY while the branch is live, AND before starting any significant new chunk of work on the branch — MANY SMALL merges, never one big deferred end-of-branch merge (small frequent merges minimise the divergence + conflict surface). (3) **HOW (safe-by-construction)** — MERGE, NEVER rebase a shared/tagged/pushed branch (no commit-loss, no history rewrite — §9/§11.4.113); FETCH-FIRST every remote before the merge (§11.4.37/§11.4.71 — remote state is unknowable without a fetch, §11.4.6); take a §9.2 hardlinked pre-op backup before any LARGE / conflict-heavy / risky merge; resolve every conflict CAREFULLY — ZERO conflict markers committed, ZERO file silently dropped, union of both sides preserved (§11.4.41 step 3 / §9 no-loss); integrate ONLY when the branch is QUIESCENT (§11.4.84 — no in-flight mutation gate, no unaccounted uncommitted work); prefer running the merge + its verification in the BACKGROUND parallel to the main stream so it never stalls other work (§11.4.88/§11.4.89/§11.4.103); NEVER force-push the merged result (§11.4.113 — the merged branch fast-forwards its own mirrors). (4) **CONSISTENCY + ANTI-BLUFF** — after EVERY merge the branch MUST be PROVEN consistent, never ASSUMED: a pre-build smoke gate runs post-merge (§11.4.42 step-3 smoke) and is GREEN, a `git grep` for conflict markers (`<<<<<<< / ===== / >>>>>>>`) returns EMPTY, and no commit present pre-merge on either side is lost — captured evidence per §11.4.5/§11.4.69; a merge claimed “clean” without the post-merge smoke + marker-scan + no-lost-commit check is a §11.4/§11.4.1 bluff at the integration layer. (5) **HONEST BOUNDARY (§11.4.6)** — regular trunk-integration keeps the branch MERGEABLE + low-conflict;

it does NOT by itself prove the branch's own work is correct (each change still crosses §11.4.108 runtime-signature + §11.4.40 retest), and it is NOT the approval-gated back-merge TO trunk (that stays §11.4.167(I)/§11.4.40/§11.4.41 no-merge-until-approved). Classification: universal (§11.4.17) — the consuming project supplies its canonical-trunk ref name, cadence unit, post-merge smoke-gate command, and per-track worktree/CoW mechanism per §11.4.35. Composes §9/§9.2/§11.4.6/§11.4.17/§11.4.37/§11.4.41/§11.4.42/§11.4.71/§11.4.84/§11.4.88/§11.4.89/§11.4.103/§11.4.113/§11.4.167/§11.4.176/§11.4.178/§11.4.179/§11.4.181. Propagation gate `CM-COVENANT-114-188-PROPAGATION` (literal `11.4.188`) + recommended gate `CM-REGULAR-MAIN-MERGE-INTO-FEATURE` (every live feature branch / track has integrated the canonical trunk within the declared cadence; its last integration is a real `git merge` not a rebase; the post-merge smoke + zero-conflict-marker + no-lost-commit checks passed) + paired §1.1 mutation (let a live branch go stale past the cadence, OR record a rebase of a shared branch in place of a merge, OR commit a conflict marker → the gate FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.188. Non-compliance is a release blocker. No escape hatch — no `--skip-main-merge`, `--merge-only-at-end`, `--rebase-shared-branch`, `--defer-trunk-sync`, `--no-post-merge-smoke`, `--force-push-after-merge`, `--big-items-only-trunk-sync` flag.

§11.4.189 — Most-reopened cases get extra-depth live-testing scrutiny FIRST (User mandate, 2026-07-10). Verbatim operator mandate: “All live testing MUST pay extra attention to cases which have been reopened numerous times and those cases in depth retested, investigated, fully validated and verified with real physical results and no bluff of any kind anywhere!” All LIVE TESTING (on-device / on-target runtime validation, deploy-and-observe) MUST give EXTRA-DEPTH retest + in-depth investigation + FULL validation/verification with REAL PHYSICAL captured evidence (§11.4.5/§11.4.69/§11.4.107 — captured audio/video/sysfs/dumpsys/sink-side/runtime-signature) and NO bluff of any kind ANYWHERE to the cases that have been REOPENED THE MOST TIMES (highest §11.4.55 reopens-count) — the empirically-most-fragile set gets the DEEPEST live scrutiny, and it gets it FIRST. (1) **WHAT** — for the highest-reopens-count cases, live testing is NOT a single pass: it re-runs the case at EXTRA DEPTH (more iterations per §11.4.50 deterministic consistency, wider edge-case + regression coverage), RE-INVESTIGATES the underlying defect per §11.4.102 systematic-debugging + §11.4.146 reproduce-first, and CONFIRMS with rock-solid captured physical evidence per §11.4.123 — never

metadata-only / config-only / absence-of-error / grep-without-runtime (§11.4/§11.4.1), never a false result. (2) **ORDER** — the most-reopened set is scrutinised FIRST, ahead of the rest of the live-test suite (the §11.4.132 risk-descending discipline, with most-reopened the load-bearing key). (3) **KEY** — priority is keyed off the §11.4.55 reopens-count (the consuming project supplies its reopens-count source per §11.4.35, e.g. the §11.4.93 workable-items DB `reopens_count`); a high reopens-count is the strongest empirical fragility signal (§11.4.132(d)). (4) **PROOF** — each most-reopened case's live PASS cites its captured-evidence artefact path (§11.4.69/§11.4.107) on a CLEAN/fresh deployment (§11.4.108 runtime-signature), with the §11.4.115 RED → GREEN polarity flip captured where a permanent regression guard exists (§11.4.135). STRENGTHENS/REFINES §11.4.132 (risk-ordered validation factor (d) most-reopened — §11.4.189 promotes most-reopened into its OWN EXTRA-DEPTH mandate at the LIVE-TEST layer, not merely one ranking factor among four) + §11.4.55 (reopens-count as the priority key) + §11.4.129/§11.4.130 (validate-the-just-fixed / highest-risk-first after redeploy). Classification: universal (§11.4.17) — references no project-specific hardware; the consuming project supplies its reopens-count source per §11.4.35. Composes §11.4.5/§11.4.6/§11.4.55/§11.4.69/§11.4.107/§11.4.108/§11.4.115/§11.4.129/§11.4.130/§11.4.132/§11.4.135/§11.4.146. Propagation gate `CM-COVENANT-114-189-PROPAGATION` (literal `11.4.189`) + recommended gate `CM-MOST-REOPENED-EXTRA-DEPTH-LIVE-TEST` (the live-test suite orders + extra-depth-scrutinises the highest-reopens-count cases FIRST, each with captured physical evidence on a clean deployment) + paired §1.1 mutation (drop the most-reopened extra-depth pass, OR run it AFTER the rest of the suite, OR PASS a most-reopened case on metadata-only evidence → the gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.189. Non-compliance is a release blocker. No escape hatch — no `--skip-most-reopened-depth` , `--reopened-same-as-any` , `--live-test-any-order` , `--metadata-pass-for-reopened-OK` , `--defer-reopened-scrutiny` flag.

§11.4.190 — Website engineering-quality mandate: every project website MUST be fully responsive + completely SEO-optimized + uniquely OpenDesign-authored + bleeding-edge enterprise-quality, each PROVEN with captured evidence (User mandate, 2026-07-10). Verbatim operator mandate: “every website we create inside our projects MUST BE fully responsive - working on all browsers, all operating systems, all platforms and all devices and screen

sizes! Every Website MUST BE completely SEO optimized! Layouts and templates done MUST BE unique and everything fully designed from OpenDesign! Make sure this applies to our project Website as well! Everything must be bleeding-edge enterprise quality, stunning visual look and feel, innovative, creative and unique! There MUST BE no false or faulty result or bluff of any kind anywhere!" Every website / web-UI surface a project ships — INCLUDING this project's OWN website — MUST satisfy AND PROVE with captured physical evidence (§11.4.5/§11.4.69/ §11.4.107/§11.4.170 — never a bare claim, §11.4.6) ALL of: **(A) FULL RESPONSIVENESS** — correct, usable rendering across ALL major browser engines (Chromium/Firefox/WebKit), all operating systems, all device classes (desktop/laptop/tablet/phone/large-display) and the full range of screen sizes/ orientations/pixel-densities; NO broken layout / overflow / clipping / overlap / off-screen or unusable control at ANY breakpoint. **(B) COMPLETE SEO OPTIMIZATION** — semantic HTML + correct document outline; per-page unique `<title>` + meta description; Open Graph + Twitter cards; canonical URLs; structured data (schema.org / JSON-LD); `robots` + XML sitemap; crawlable / indexable markup; WCAG AA accessibility + Core Web Vitals performance (both feed ranking). **(C) UNIQUE OpenDesign-AUTHORED templates/layouts** — EVERY layout + template designed FROM the OpenDesign design-token system (§11.4.162), NOT generic / off-the-shelf / templated; visually unique, innovative, creative, brand-distinct. **(D) BLEEDING-EDGE ENTERPRISE VISUAL QUALITY** — stunning, polished, professional look-and-feel in BOTH light AND dark themes (§11.4.162 tokens). **(E) ANTI-BLUFF PROOF** — NONE of (A)–(D) may be CLAIMED without captured physical evidence: responsiveness PROVEN by device-independent host-rendered screenshots across the real breakpoint × engine matrix + an OCR/vision layout oracle (no overlap / label-over-label / clip / overflow / off-screen); SEO PROVEN by an automated audit (Lighthouse SEO + structured-data validation + meta / OG / sitemap / robots presence) meeting a defined score floor; visual quality PROVEN by the §11.4.170 device-independent host-rendered pixel proof per screen × state × {light,dark}; uniqueness PROVEN by OpenDesign-token provenance. A website reported responsive / SEO-optimized / unique / enterprise-quality WITHOUT these captured proofs is a §11.4 PASS-bluff at the web-UI layer. §11.4.190 is the strict WEB-quality expansion of §11.4.162 (OpenDesign — it applies §11.4.162's design-token mandate to the full responsive + SEO + enterprise-quality web surface) and binds §11.4.170's device-independent host-rendered pixel proof as the correctness oracle; sibling of §11.4.168 (exported-document independent visual validation) at the web-UI layer. Classification:

universal (§11.4.17) — references no project-specific hardware; the consuming project supplies its concrete breakpoint × engine matrix, SEO score-floor, and OpenDesign token set per §11.4.35. Composes §11.4.5/§11.4.6/§11.4.69/§11.4.107/§11.4.162/§11.4.168/§11.4.170. Propagation gate `CM-COVENANT-114-190-PROPAGATION` (literal `11.4.190`) + recommended gates `CM-WEBSITE-RESPONSIVE-PROVEN` (host-rendered screenshots across the breakpoint × engine matrix + layout oracle — no overlap/clip/overflow/off-screen) / `CM-WEBSITE-SEO-OPTIMIZED` (automated SEO audit meets the score floor + structured-data / meta / OG / sitemap / robots present) / `CM-WEBSITE-OPENDSIGN-UNIQUE` (every layout/template's OpenDesign-token provenance, no generic/off-the-shelf template) + paired §1.1 mutation (report a website responsive/SEO/unique/enterprise-quality with NO captured proof, OR ship a generic non-OpenDesign template, OR skip the light/dark host-rendered pixel proof → the recommended gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.190. Non-compliance is a release blocker. No escape hatch — no `--skip-responsive-proof`, `--seo-optional`, `--generic-template-OK`, `--visual-quality-later`, `--claim-without-render` flag.

§11.4.191 — Work-to-track/branch binding enforcement: a logic-group's work can NEVER be committed OR dispatched onto the wrong track/branch (research-derived, 2026-07-10). Compact summary: every logical group of workable items binds to EXACTLY ONE canonical (`track`, `branch`) recorded ONCE in the single source of truth (the §11.4.93 workable-items DB `logic_groups.destination` = authoritative branch + a `canonical_track` column + a `group_paths` file-scope manifest), and NO change belonging to a group may LAND on, nor be DISPATCHED to, a track/branch other than that group's canonical one — the file-scope generalisation of §11.4.181/§11.4.182 from branch-NAME consistency to work-PLACEMENT consistency (forensic FACT 2026-07-10: `srcpy/Remote-app` work — group `mistiq-vader`, canonical `feature/mistiq-vader` /T2 — was committed onto `main` /T1; §11.4.181's `guard-branch-consistency.sh` never fired because committing onto an existing branch is not a feature-branch CREATE, and no gate checked which FILES land on which BRANCH). (1) **Single source of truth** — the `logic_group` → (`branch`, `track`, `path-globs`) binding is recorded ONCE in the DB and LOOKED UP, never re-invented (a commit/dispatch placing a group's file-scope on a non-canonical branch/track is a §11.4.6 no-guessing violation). (2) **Preventive hook**

(§11.4.109/§11.4.177) — a PreToolUse guard `guard-work-track-binding.sh` (inherited BY REFERENCE, never copied) BLOCKS (exit 2) a git-commit-class Bash call whose to-be-committed file set contains a path owned by a group whose canonical branch/track $\neq$ the current checkout, AND BLOCKS an Agent/Task dispatch whose §11.4.182 label track/branch $\neq$ the canonical (track, branch) of the ticket's group; fail-closed on unreadable registry, `# guardrails:allow <reason>` audited escape. (3) **Detective gate (§11.4.110/§11.4.108)** — `CM-WORK-TRACK-BINDING-ENFORCED` asserts, on the current branch's own added commits (`merge-base(HEAD,main)..HEAD`), that every `group_paths`-owned file's group destination == current branch (and `canonical_track` == current track when pinned) — the actual enforcement boundary regardless of how the commit was made, the runtime signature being FAIL-on-wrong-placement / PASS-on-canonical-placement on a clean checkout. (4) **Risk-free reconciliation (§9.2/§11.4.113/§11.4.84)** — a mis-placed change is reconciled by MERGING its commits onto the group's canonical branch (union, no commit lost, no `-s ours /reset`), ff-only push to all upstreams (NEVER force-push), then removing it from the wrong branch only after the merge is confirmed on all mirrors. Classification: universal (§11.4.17) — references no project hardware; the consuming project supplies its `(track, branch)` map, `group_paths` seed, and commit-wrapper wiring per §11.4.35. Composes

§11.4.6/.28/.29/.35/.58/.75/.84/.93/.95/.109/.110/.113/.119/.167/.176/.177/.178/.181/.182/

§9.2. Propagation gate `CM-COVENANT-114-191-PROPAGATION` (literal `11.4.191`) + recommended gate `CM-WORK-TRACK-BINDING-ENFORCED` + paired §1.1 mutation.

Canonical authority: constitution submodule `Constitution.md` §11.4.191. Non-compliance is a release blocker. No escape hatch — no `--allow-cross-track-commit`, `--skip-work-binding`, `--commit-any-branch-OK`, `--dispatch-any-track-OK`, `--no-file-scope-check` flag.

§11.4.192 — Continuous multi-track auto-backfill: a FREE track MUST be IMMEDIATELY re-assigned its next-highest-priority domain work-set, never left idle while its domain has actionable items (User mandate, 2026-07-11).

Verbatim operator mandate: “once a Track is free, we shall assign next set of workable items to it ... full continuous multi-track development ... no stopping ... this is MANDATORY”. Compact summary: the moment any track's worker COMPLETES its current work-set (its item(s) closed with anti-bluff evidence, or merged), that FREE track MUST be IMMEDIATELY assigned the NEXT-highest-priority actionable work-set from its OWN domain (its §11.4.191 `logic_group` /

canonical branch) — a track is NEVER left idle while its domain still holds actionable workable items, and the conductor/ruler CONTINUOUSLY re-assigns with NO stopping (full continuous multi-track development); PROMOTES the §11.4.103(B) ≥3-parallel-background-streams-with-auto-backfill invariant into the multi-track (§11.4.176/§11.4.178/§11.4.179/§11.4.187) layer (§11.4.103(B) backfills a STREAM the moment it is done; §11.4.192 backfills a TRACK the moment its worker is done, keyed to that track's canonical (track, branch) domain per §11.4.191). (1) **WHAT** — on any track-worker completion the conductor/ruler (§11.4.187) MUST, per that track's domain, look up the NEXT actionable work-set, CLAIM it exactly-once (§11.4.176 exactly-once claim registry, no two tracks claiming one item), and DISPATCH it as the track's new §11.4.182-labelled work stream — without waiting for a scheduled wake or an operator prompt. (2) **PRIORITY** — the next work-set is the track domain's highest-priority actionable item (§11.4.42/§11.4.72 priority order, audio-first where in scope), never an arbitrary pick. (3) **SOURCE-SIDE ADVANCES EVEN WHEN VALIDATION IS GATED** — when the next item's on-device VALIDATION is blocked on a pending build / device availability (§11.4.108 clean-target / §11.4.185 manual-QA / operator-gated flash), the track MUST still advance the item's SOURCE-SIDE work (investigation, RED-test §11.4.115, source patch, pre-build gate + paired §1.1 mutation, docs) and honestly §11.4.3/§11.4.52 SKIP-with-reason ONLY the genuinely device-gated validation portion — a build/device gate on the validation step is NOT a reason to idle the whole track. (4) **HONEST BOUNDARY (§11.4.6)** — a track idles ONLY when its domain GENUINELY has zero remaining actionable workable items OR every remaining item is externally blocked with no source-side advancement possible (§11.4.94(A)/§11.4.97/§11.4.101 idle-only-when-blocked closed set); “the merges are done” / “the current item is done” / “nothing visible right now” is NEVER a valid idle reason while the domain still holds progressable items (a §11.4.94+§11.4.97 violation); an unavoidable per-item block PARKS one item, it does NOT stop the track (§11.4.101 — the track backfills its next non-blocked domain item). Classification: universal (§11.4.17) — references no project hardware; the consuming project supplies its track → domain map + priority source per §11.4.35. Composes §11.4.42/.58/.72/.94/.97/.101/.103/.167/.176/.178/.179/.182/.187. Propagation gate CM-COVENANT-114-192-PROPAGATION (literal 11.4.192) + recommended gate CM-CONTINUOUS-MULTITRACK-BACKFILL (on a track-worker completion the ruler re-assigns that track's next-highest-priority domain work-set; a track with a completed worker AND remaining actionable domain items but no new dispatch →

the gate FAILs) + paired §1.1 mutation (let a free track sit idle while its domain holds an actionable item, OR treat “merges done” as a valid idle reason → the gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item).

Canonical authority: constitution submodule `Constitution.md` §11.4.192. Non-compliance is a release blocker. No escape hatch — no `--allow-idle-track`, `--skip-backfill`, `--stop-multitrack`, `--merges-done-is-idle`, `--defer-track-reassign` flag.

§11.4.193 — Anti-blind-typing mandate: every UI interaction MUST be “seen” and “understood” via OCR/vision/screenshot proof, NEVER blind-typed (User mandate, 2026-07-13). Verbatim operator mandate: “No blind typing of any kind anywhere!!!”; “Login to apps and services MUST BE done with understanding and seeing the whole UI ALWAYS!!!”; “blind typing is STRICTLY FORBIDDEN, and instead everything MUST BE seen and understood in coordination with proper models and means we have incorporated in HelixQA!!!”. Blind typing — sending keystrokes/text/input events to a UI surface WITHOUT first capturing AND understanding what is on screen (OCR / vision / uiautomator / accessibility-tree / screenshot) AND confirming the result after each interaction — is STRICTLY FORBIDDEN everywhere and always, for every agent/tool/automation/platform/context. It is LITERAL bluff: claiming “credentials entered”/“login completed” without visual proof the keystrokes landed on the intended fields/screen producing the intended result is a §11.4 PASS-bluff at the UI-interaction layer (a credential typed into the wrong field, or a WebView that silently swallowed input, = end-user feature broken while the agent reported success). (1) **SEE-BEFORE-YOU-TYPE** — before ANY `input text` / `input keyevent` / `adb shell input` / uiautomator gesture, capture the current screen state (screencap + Tesseract OCR ≥150 DPI, OR uiautomator dump grep, OR accessibility-tree dumphsys, OR HelixQA VisionEngine frame per §11.4.159/§11.4.160) AND confirm the intended target element is visible AND focused — the capture IS the §11.4.5/§11.4.69 `touch_input` evidence. (2) **VERIFY-AFTER-YOU-TYPE** — after EVERY interaction step, capture a NEW screen state + confirm the intended change (5 `input text` + 1 final screenshot = blind typing). (3) **OCR/VISION ORACLE MANDATORY** — Tesseract ≥150 DPI / uiautomator-dump grep / HelixQA VisionEngine / accessibility-tree dumphsys. (4) **WEBVIEW** — near-empty uiautomator hierarchy → OCR fallback per §11.4.117; if even OCR is infeasible (secure/DRM-blanked surface captures black) → the login is operator_attended-only, SKIP-with-reason per §11.4.3, NEVER blind-type-and-claim-success. (5)

CREDENTIAL-SAFE OCR — redact credential characters from OCR output + logs per §11.4.10; evidence = the redacted screenshot / OCR text. (6) **ANTI-BLUFF** — every login/interaction PASS cites a captured-evidence artefact containing pre-interaction state + post-interaction state + OCR/vision verification of the expected change; self-validated per §11.4.107(10) with a golden-good + golden-bad fixture pair. Classification: universal (§11.4.17). Composes §11.4 / §11.4.5 / §11.4.6 / §11.4.69 / §11.4.107 / §11.4.117 / §11.4.159 / §11.4.160 / §11.4.10 / §11.4.52 / §11.4.3. Propagation gate `CM-COVENANT-114-193-PROPAGATION` (literal `11.4.193`) + recommended gate `CM-NO-BLIND-TYPING` + paired §1.1 mutation (strip OCR verification before an `input text` → gate FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.193. Non-compliance is a release blocker. No escape hatch — no `--blind-typing-ok`, `--skip-ocr-verification`, `--assume-input-landed`, `--login-without-proof`, `--type-without-seeing` flag.

§11.4.194 — Exhaustive all-scenario, all-angle code-review + verify-against-captured-runtime-evidence mandate (User mandate, 2026-07-14). Verbatim operator mandate: “CRITICAL: Code review MUST TAKE into the accounts all scenarios, look at the problem or change from all angles, deepest possible analysis! No gaps, shortcomings or false results or bluff is allowed! Extend our root constitution Submodule and all relevant files properly so the review process is more precise and we do not have situations like the last one!” STRENGTHENS §11.4.125/§11.4.134/§11.4.142 — every code-review MUST, before GO: (1) ENUMERATE the FULL input/scenario space of the change AND the defect (every dimension/path/interacting feature, not just the reported scenario); for a MULTI-FACTOR failure mode (a product of N terms, ANY degenerate term — zero/empty/null/default-init/out-of-range — triggers the failure) enumerate + INDEPENDENTLY VERIFY the fix covers EACH factor — verifying one and assuming the rest is the exact forbidden gap (forensic FACT “the last one”: a boot-crash fix got a clean zero-finding GO that dismissed a clamp branch as an informational NIT “effectively unreachable ... always valid” WITHOUT PROVING it against the runtime input space; the target STILL crash-looped because the degenerate default-init input the review assumed impossible DID happen — a `size = A × B` product (e.g. `frame_size = channels × bytes_per_sample(format)`) with TWO independent zero-triggers, only ONE factor verified, the OTHER assumed away, the pre-fix crash evidence never cross-checked); (2) PROVE every assumption —

“can’t happen”/“unreachable”/“always valid”/“caller never sends X” MUST be PROVEN from code+spec+captured evidence, never on faith (§11.4.6); an unproven assumption that proves false is a §11.4 review-layer bluff of PASS-bluff severity; a NIT dismissing a path as unreachable MUST carry the proof of unreachability (a bare “effectively unreachable” NIT is itself a finding); (3) VERIFY the fix against ANY captured runtime evidence (crash log/tombstone/stack-trace/sink-probe/§11.4.108 runtime-signature/reproducer) — confirm the fix ACTUALLY prevents the CAPTURED failure against the exact captured conditions; a GO on a fix that then fails at runtime on a scenario present in the evidence (or one the review chose not to consider) is a bluff; (4) ALL ANGLES — correctness, full input domain (every factor/boundary/degenerate value), every downstream/interacting must-not-break feature (§11.4.133 target/hardware safety + shared-state features e.g. audio/multichannel/passthrough), the TESTS’ OWN per-dimension coverage (a test checking only one factor of a multi-factor bug is itself a gap the review catches), the §1.1 mutation per dimension; (5) HONEST BOUNDARY (§11.4.6) — “exhaustive” = the enumerated + evidence-grounded scenario space, NOT literal omniscience; an un-analysable dimension is stated as an explicit gap (UNCONFIRMED/UNKNOWN/PENDING_FORENSICS + tracked follow-up), NEVER silently assumed safe. Classification: universal (§11.4.17). Composes §11.4.125/.134/.142/.6/.108/.130/.107(10). Propagation gate CM-COVENANT-114-194-PROPAGATION (literal 11.4.194) + recommended gate CM-EXHAUSTIVE-REVIEW-ALL-SCENARIOS + paired §1.1 mutation. **Canonical authority:** constitution submodule Constitution.md §11.4.194. Non-compliance is a release blocker. No escape hatch — no `--partial-scenario-review` , `--assume-unreachable-OK` , `--skip-captured-evidence-crosscheck` , `--single-factor-review-suffices` , `--dismiss-branch-without-proof` flag.

§11.4.195 — Branch-structure governance: taxonomy + merge-after-live-QA + flavor/product non-merge (User mandate, 2026-07-14). Verbatim operator rules: branch taxonomy `feat/<NAME_OF_FEATURE>` / `product/<NAME_OF_PRODUCT>` / `flavor/<NAME_OF_FLAVOR>` ; tracks on FEATURE branches or `main` MUST — after full confirmation + validation + verification + full LIVE MANUAL TESTING (§11.4.185) — MERGE to `main` COMPLETELY (all submodules + main repo), ff-only, NEVER force-push (§11.4.113), no commit lost (§9.2); tracks with an entirely-new FLAVOR/PRODUCT do ALL work on a `flavor/` AND/OR `product/` branch (both may exist) which does NOT merge to `main` the same way; commit + push ALL branches to ALL upstreams (§2.1), ff-only. Every controlled branch belongs to

one of THREE canonical taxonomy classes, each with a distinct integration lifecycle: **(A) Taxonomy** (extends §11.4.181/§11.4.167(D)) — `feat/<slug>` = FEATURE (merges to `main` after live-manual-QA, clause B); `product/<slug>` = PRODUCT line + `flavor/<slug>` = FLAVOR line (do NOT merge to `main` the standard way, clause C); `<slug>` lowercase snake/kebab (§11.4.29); the canonical branch NAME is minted ONCE + used IDENTICALLY across main repo + every touched owned submodule (§11.4.181), recorded once in the §11.4.176 claim registry / §11.4.93 `logic_group` → `branch` map + LOOKED UP never re-invented (§11.4.6); the prefix STRINGS are operator-canonical, a project already on a different feature-prefix (`feature/...`) is compliant so long as it is the ONE canonical name per §11.4.181, reconciled by the §11.4.181 risk-free MERGE-onto-canonical path (ff-only, no commit lost, never a rename that drops commits). **(B) FEATURE/main** → **merge-after-live-QA** — before merging to `main` satisfy IN ORDER automated four-layer (§11.4.4(b)) + §11.4.40 full-suite retest THEN full LIVE MANUAL TESTING by the QA team (§11.4.185 FINAL human gate, automated GREEN necessary NOT sufficient); only after §11.4.185 may the merge proceed + it MUST be COMPLETE (main repo + every touched owned submodule) ff-only onto latest `main` (§11.4.113), NEVER force-push, NO commit lost / NO history rewrite (§9.2/§11.4.41 step 3), pushed to ALL upstreams (§2.1); binds the §11.4.185 gate as an explicit precondition of the §11.4.167(I)/§11.4.188 approval-gated back-merge. **(C) FLAVOR/PRODUCT lines** → **do-not-merge-to-main-the-same-way** — an entirely-new flavor/product line does ALL work on a `product/<NAME>` AND/OR `flavor/<NAME>` branch (both may coexist), a divergent line with its own release lifecycle that does NOT flow back to `main` via clause B; STILL obeys §11.4.113 (no force-push) / §9.2 (no commit lost) / §2.1 (push all branches all upstreams ff-only) / §11.4.188 (MAY merge `main` INTO itself to stay close — trunk-into-line permitted; only line-into-trunk is gated/withheld). **(D) Push discipline** (§2.1/§11.4.113) — commit + push ALL branches (feature/product/flavor) to ALL upstreams ff-only, NEVER force-push; new-branch creation is inherently ff (§9.2) + non-disruptive from an existing HEAD, permitted in a dirty tree under §11.4.84 provided no commit is in-flight. Honest boundary §11.4.6 — governs branch STRUCTURE + integration DIRECTION, does NOT prove any branch's work correct (still §11.4.108 runtime-signature + §11.4.40 retest + §11.4.185 manual-QA). Classification: universal (§11.4.17) — strict generalisation of §11.4.181 + §11.4.167(D)/(I) + §11.4.188 + §11.4.185. Composes §2.1/§9/§9.2/§11.4.6/§11.4.17/§11.4.29/§11.4.35/§11.4.40/§11.4.41/§11.4.84/§11.4.93/§11.4.113/

§11.4.167(D)/(I)/§11.4.176/§11.4.178/§11.4.181/§11.4.185/§11.4.188/§11.4.191.
 Propagation gate `CM-COVENANT-114-195-PROPAGATION` (literal `11.4.195`) +
 recommended gate `CM-BRANCH-TAXONOMY-CLASS` (every controlled branch prefix
 $\in$ `feat/` | `product/` | `flavor/` | reserved `main` / `master` ; a `product/` |
`flavor/` branch is NOT the source of a merge onto `main` ; a `feat/` | `main`
 back-merge cites a §11.4.185 manual-QA confirmation) + paired §1.1 mutation (a
`product/*` merged onto `main` , OR a `feat/*` back-merge with no §11.4.185
 citation → gate FAILs; strip the literal → propagation gate FAILs; gate-code =
 separate work item). **Canonical authority:** constitution submodule
`Constitution.md` §11.4.195. Non-compliance is a release blocker. No escape
 hatch — no `--any-branch-prefix-OK` , `--merge-product-to-main` , `--skip-`
`live-qa-before-merge` , `--force-push-branch` ,
`--partial-submodule-merge` flag.

§11.4.196 — Native-alias-first priority + per-alias real-signal limit/subscription tracking + auto-rebind-on-recovery + resource-detection-by-real-identity-not-substring-match (User mandate, 2026-07-14). Verbatim operator mandate: use ALL operational claude-NATIVE aliases FIRST — as long as ≥ 1 native alias is operational (NOT session-rate-limited, NOT weekly-limit-reached, NOT expired-subscription) — before ANY provider alias; track per alias the session rate-limit / weekly-limit / subscription expiry+renewal from REAL signals (never faked, §11.4.6) + record when each becomes operational AGAIN; the moment a higher-priority native recovers, rebind tracks to it; and fix the host-budget guard footgun where a `pgrep -f substring-match false-REFUSED` every worker spawn. The multi-track alias-binding + rate-limit/subscription-tracking layer of the §11.4.187 ruler orchestration MUST implement ALL of: **(A) NATIVE-ALIAS-FIRST PRIORITY** — every OPERATIONAL native (`claudeN`) alias is selected before ANY provider, UNCONDITIONALLY, while ≥ 1 native is operational; the guarantee is STRUCTURAL — a native/provider CLASS partition scanned native-class-completely-before-provider-class (§11.4.111 resolve-by-CLASS-not-by-file-position), so a provider NEVER wins while a non-cooled native exists regardless of roster order; within a class the operator-mandated order (native equal-capability, then providers `deepseek` → `xiaomi` → `opencode` → `kimi` → ...) is authoritative decision-DATA (§11.4.28) never re-invented (§11.4.6). **(B) PER-ALIAS REAL-SIGNAL LIMIT + SUBSCRIPTION TRACKING** — “operational?” per alias from REAL captured signals, NEVER faked/guessed (§11.4.6): the §11.4.187 captured 429 signature (`isApiErrorMessage` + `apiErrorStatus:429` / `api_error_status:429`) sub-

classified into reason-CLASS closed set {session | weekly | subscription} from real markers (reset/weekly-limit/session-limit); each limited alias records its CLASS + an operational-again EPOCH with a DIFFERENT window per class — session (self-heals at reset, short), weekly (~until the weekly reset, long), subscription (INDEFINITE — operational-again UNKNOWN until a REAL renewal signal, parked with a far-future sentinel; §11.4.6 NEVER guesses a renewal date, an operator/config supplies it). A limited alias is NEVER auto-selected until its epoch passes (session/weekly) or a real renewal/recovery clears it (subscription).

(C) AUTO-USE-ON-RECOVERY — the moment a higher-priority native recovers (window elapsed / renewal marked), rebind tracks UPWARD to it (a promote -class op keyed off a comparable priority-rank: native class before provider, then config order), PRESERVING each track's worktree AND device leases (leases key on the STABLE track id — the alias changes, not the track; §11.4.119/§11.4.187); idempotent. **(D) RESOURCE-DETECTION BY REAL IDENTITY, NOT SUBSTRING-MATCH (RB-02 guard fix)** — a “heavy work in flight” guard MUST identify a real resource by its ACTUAL process command line read from the OS (/proc/<pid>/cmdline), NEVER a bare pgrep -f REGEX substring match that false-matches a process merely MENTIONING a token — a pattern-CARRIER (a process quoting the whole alternation pattern; a claude / worker whose prompt embeds the token; a launcher/monitor/test), a same-tool ENGINE process, or the guard's own process/parent (the §12.12 self/sibling pgrep footgun — forensic FACT 2026-07-13: the guard false-REFUSED EVERY §11.4.187 spawn on a host with ZERO real builds because a live claude worker's prompt quoted the pgrep pattern); the guard re-reads each matched PID's real cmdline + excludes those carrier/engine/worker/self classes; an UNREADABLE/ vanished cmdline is conservatively counted as the real resource (REFUSE — safe reversible default, §11.4.101/§12.8), so a genuine build is still caught. **(E) ANTI-BLUFF (§11.4/§11.4.108)** — four-layer coverage each (pre-build gate + on-host test + paired §1.1 mutation + captured evidence §11.4.69) with §11.4.115 RED-polarity + §11.4.50 determinism; honest boundary §11.4.6 — mandates the MECHANISM + REAL-signal derivation of “operational”, genuine live per-subscription quota-isolation proof still needs live tokens (operator-scheduled acceptance window), never conflated with a mock/hermetic-contract GREEN. Classification: universal (§11.4.17). Composes §11.4.187/.182/.176/.111/.101/.103/.6/.108/.115/.119/§12.12/§12.6/§12.8. Propagation gate CM-COVENANT-114-196-PROPAGATION (literal 11.4.196) + recommended gate CM-ALIAS-PRIORITY-LIMIT-TRACKING (native-first CLASS

partition; per-alias reason-class + operational-again tracking; auto-rebind-on-recovery; build-guard resolves by real cmdline not bare substring) + paired §1.1 mutation (swap native/provider priority bases so a provider outranks a native → recommended gate FAILs; flip the guard's strict-cmdline-filter to the pre-fix bare `pgrep` → its RB-02 test FAILs; strip the literal → propagation gate FAILs; gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.196. Non-compliance is a release blocker. No escape hatch — no `--provider-before-operational-native`, `--fake-limit-state`, `--guess-renewal-date`, `--no-rebind-on-recovery`, `--substring-match-build-guard`, `--skip-real-signal-classification` flag.

§11.4.197 — Research / kicked-off-work completion mandate: every research effort MUST be driven to full, wired, verified completion or explicitly closed — never left un-wired in the backlog (User mandate, 2026-07-15). Verbatim operator mandate: “all research work we start MUST BE fully finished, completely confirmed, validated and verified! It MUST NOT happen for it to stay at the backlog uncompleted and un-wired! Loss of requirements, failure to incorporate them is FORBIDDEN!!! Everything MUST BE fully completed! No exceptions!” Every research effort / kicked-off improvement / design doc / spike / prototype / PoC / partially-landed feature a project STARTS MUST be driven to FULL completion OR explicitly evidence-backed CLOSED — there is no third “sitting un-wired in the backlog” state; “we researched/designed/started it” is NEVER terminal, and a `docs/research/*` doc / design doc / spike / half-wired feature that is neither COMPLETED-and-wired nor explicitly-closed is a §11.4.197 violation of §11.4 PASS-bluff severity at the requirements-integrity layer (the requirement was accepted, work began, then it silently evaporated); loss of requirements + failure to incorporate started work are FORBIDDEN. **COMPLETED** = (1) implemented + (2) WIRED (integrated AND used by default, NOT merely present behind a dead flag / uncalled function / unreferenced module / never-selected path — mere-presence-without-wiring is the §11.4.124 dead-code + §11.4.108 layer-2/3 SOURCE → ARTIFACT → RUNTIME gap) + (3) confirmed/validated/verified with captured physical evidence §11.4.5/§11.4.69/§11.4.107 on a clean deployment §11.4.108 runtime-signature (never metadata-only/config-only/absence-of-error/grep-without-runtime §11.4/§11.4.1) + (4) tracked to a terminal Status §11.4.15/§11.4.33. **Explicitly CLOSED** = deliberately terminated with a documented evidence-backed reason from the closed vocabulary — §11.4.90 Obsolete / §11.4.112 structurally-impossible / §11.4.122 operator-confirmed drop — NEVER

silent abandonment (“no time”/“lower priority” = §11.4.21 Operator-blocked or a still-open tracked item, NOT a close). **Mechanical tracker binding (§11.4.93):** every research effort / design doc / kicked-off improvement MUST map to ≥1 tracked workable item (ATM-NNN §11.4.54) whose completion the loop enforces; a research doc with NO implementing item, or a tracked item stalled un-wired past the cadence, is a violation; the autonomous loop (§11.4.87/.94/.97/.103/.126/.192) MUST treat every un-wired research effort as an actionable queue item — driving it to COMPLETED or explicitly CLOSED is exactly the “keep working until the queue is genuinely empty” contract. Honest boundary §11.4.6 — “fully completed” is bounded by wired+verified OR evidence-backed-closed, NOT “every speculative idea ships”; every STARTED effort reaches a documented terminal state, proven never assumed. Classification: universal (§11.4.17). Composes §11.4.8/.87/.94/.97/.103/.118/.123/.124/.150/.192/.90/.112/.122/.21/.93/.54/.108/.15/.33/.126. Propagation gate `CM-COVENANT-114-197-PROPAGATION` (literal `11.4.197`) + recommended gate `CM-RESEARCH-COMPLETION-TRACKED` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.197. Non-compliance is a release blocker. No escape hatch — no `--research-may-stay-unwired`, `--skip-completion-tracking`, `--abandon-without-close`, `--present-but-unwired-is-done`, `--backlog-research-OK`, `--lose-requirement-OK` flag.

§11.4.198 — Default working mechanisms: multi-alias native-first orchestration AND heavy token-optimization are ALWAYS-ON defaults for single- and multi-track work (User mandate, 2026-07-15). Operator mandate: multiple claude-toolkit aliases (native-first) AND heavy token-optimization MUST be the DEFAULT mechanisms, always used with single- and multi-track development and with single or multiple working agents. Two proven mechanisms promoted to STANDING always-on defaults, engaged automatically for EVERY configuration (single- OR multiple-agent, single- OR multi-track), never a per-request opt-in: **(A) MULTI-ALIAS NATIVE-FIRST ORCHESTRATION** — the claude-toolkit multi-alias mechanism (every OPERATIONAL claude-native alias before ANY provider per §11.4.196 native-alias-first priority + per-alias real-signal limit/subscription tracking + auto-rebind-on-recovery, driven by the §11.4.187 automatic multi-track ruler) MUST be the default execution substrate for ALL work, NOT only the multi-track case; a SINGLE-track/SINGLE-agent session STILL binds its worker through the same native-first alias-priority mechanism (best operational native + rebind-on-recovery), a MULTI-track/multiple-agent config uses it across every track (§11.4.176

exactly-once claim + §11.4.182 track+branch+alias label); running on an arbitrary/hardcoded/provider-first alias when an operational native exists, or bypassing the mechanism because “it is just one agent”, is a violation. **(B) HEAVY TOKEN-OPTIMIZATION** — the §11.4.141 token-efficiency discipline (compact-anchor layout with full-one-hop-away, targeted reads over whole-file dumps, subagent context isolation §11.4.20/§11.4.70, evidence-hash references not evidence dumps) **MUST** apply **ALWAYS** — every read/dispatch/doc-edit — for single- AND multi-track, single OR multiple agents; token-heavy patterns (re-reading whole large files when a targeted slice suffices, dumping full evidence into context, monolithic non-subagent execution of parallelisable work) when a token-efficient path exists are violations. Honest boundary §11.4.6 — “default” = auto-engaged without an operator request; it does NOT weaken host-safety, the multi-alias/multi-agent fan-out stays under the §12.6 60% memory ceiling + §12.12 thread-headroom + §11.4.58 parallel-agent cap (contending/redundant agents are NOT “the default”, §11.4.183 maximum-USEFUL-throughput); §11.4.198 BINDS the two mechanisms as always-on, it does NOT redefine them (mechanics live in §11.4.196/§11.4.187/§11.4.141). Classification: universal (§11.4.17). Composes §11.4.141/.187/.196/.182/.176/.183/.103/.58/.20/.70/.126/§12.6/§12.12. Propagation gate `CM-COVENANT-114-198-PROPAGATION` (literal `11.4.198`) + recommended gate `CM-DEFAULT-MECHANISMS-ALWAYS-ON` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.198. Non-compliance is a release blocker. No escape hatch — no `--single-agent-skips-alias-priority`, `--provider-first-default`, `--token-optimization-optional`, `--heavy-tokens-OK`, `--not-default-mechanism` flag.

§11.4.199 — Exact-reproduction-sequence mandate: when a working reproduction exists, the investigation MUST use ITS exact sequence — a deviating repro proves NOTHING (research-derived, 2026-07-15). Forensic FACT (2026-07-15): an investigation concluded a defensive “cover” mechanism “never engages” (and the defect it guards was therefore unfixed) — the conclusion was **WRONG** because the hand-rolled reproduction drove the target with a **DIFFERENT** input than the existing test (a HOME key where the test sent a media-pause key), so it never reached the PRECONDITION the mechanism keys off; re-driven with the TEST’S EXACT sequence the mechanism provably **ENGAGED**. Every investigation / debugging pass / validation that reproduces a defect or exercises a code path **MUST**: (1) **use the existing reproduction’s EXACT sequence** when a working reproduction already exists (a test/harness/recorded

sequence that demonstrably reaches the defect) — same inputs, same order, same timing, same preconditions, same topology (§11.4.3); a hand-rolled approximation is NOT a substitute; (2) **a deviating variant MUST FIRST prove it reaches the precondition** with captured evidence before ANY conclusion is drawn from it; (3) **a repro that misses the precondition proves NOTHING** — it is NOT evidence of absence (§11.4.6), and concluding “defect present”/“defect absent”/“the mechanism never engages” from it is a §11.4/§11.4.1 bluff at the INVESTIGATION layer (a FAIL-bluff when it wrongly condemns working code, a PASS-bluff when it wrongly clears broken code); (4) **record the exact driving sequence as captured evidence** so it is replayable + auditable, every deviation explicit + justified + precondition-proven. STRENGTHENS §11.4.102 (Phase-2 reproduce) / §11.4.146 (reproduce-first) / §11.4.115 (RED on the broken artifact) — they mandate THAT you reproduce, §11.4.199 mandates HOW: with the sequence PROVEN to reach the defect. Classification: universal (§11.4.17). Composes §11.4.1/.3/.5/.6/.7/.102/.107/.115/.123/.135/.146. Propagation gate CM-COVENANT-114-199-PROPAGATION (literal 11.4.199) + recommended gate CM-EXACT-REPRO-SEQUENCE + paired §1.1 mutation. **Canonical authority:** constitution submodule Constitution.md §11.4.199. Non-compliance is a release blocker. No escape hatch — no --approximate-repro-OK, --skip-precondition-proof, --hand-rolled-repro-suffices, --absence-proves-absence, --deviating-sequence-conclusion-OK flag.

§11.4.200 — Non-targetable deploy/flash tooling MUST isolate exactly ONE eligible target and VERIFY-AFTER-WRITE on the INTENDED target (research-derived, 2026-07-15). Forensic FACT (2026-07-15): a flashing tool with no device-selection argument AUTO-SELECTED a device from the bus, silently wrote the freshly-built image to the WRONG board (a sibling still attached), printed SUCCESS, and the INTENDED board kept its OLD build — every downstream test then validated a target that had NEVER received the artifact (the §11.4.108 ARTIFACT → RUNTIME layer silently broken while every tool reported green). Whenever a deploy / flash / install / write tool CANNOT address a specific target, the procedure MUST: (1) **targetability is a PRECONDITION** — EITHER pass an explicit STABLE selector (serial/UUID/stable id per §11.4.111, never an enumeration ordinal) OR GUARANTEE exactly ONE eligible target is present at the moment of the write (siblings disconnected/powered-down/excluded), the enumerated eligible-target set captured as evidence (size == 1); (2) **the tool’s own success message is NOT proof** — an “ok”/exit-0 proves only that bytes reached

the device the TOOL selected, NEVER the device the OPERATOR intended (§11.4.6); (3) **VERIFY-AFTER-WRITE on the intended target** — read back the deployed artifact's IDENTITY (build id / version / checksum / fingerprint) FROM the intended target and assert it equals the artifact just written; a mismatch or the OLD identity is a FAIL, never a warning (the §11.4.108 ARTIFACT → RUNTIME assertion applied to the deployment step); (4) **no validation on an unverified target** — any test run against a target whose deployed-artifact identity was NOT verified is INVALID and any PASS is a §11.4 PASS-bluff. Honest boundary §11.4.6 — isolation + read-back prove the INTENDED target holds the INTENDED artifact, NOT that the artifact is correct (§11.4.108/.40/.130). Classification: universal (§11.4.17). Composes §11.4.6/.46/.108/.111/.119/.130/.133/.139. Propagation gate `CM-COVENANT-114-200-PROPAGATION` (literal `11.4.200`) + recommended gate `CM-DEPLOY-TARGET-ISOLATED-AND-VERIFIED` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.200. Non-compliance is a release blocker. No escape hatch — no `--auto-select-target-OK`, `--trust-tool-success-message`, `--skip-post-write-verify`, `--validate-unverified-target`, `--multiple-targets-on-bus-OK` flag.

§11.4.201 — Every guard/gate MUST assert the REAL condition: a false-positive refusal is a FAIL-bluff, a false-negative pass is a PASS-bluff (research-derived, 2026-07-15). Forensic FACT (2026-07-15, two independent instances in one cycle): (a) a host-budget guard used a bare `pgrep -f '<pattern>'` and REFUSED every worker spawn on a host with ZERO real builds because a live worker's own command line merely QUOTED the pattern (the §11.4.196(D)/§12.12 carrier footgun); (b) a host-safety pre-flight refused a build on a “swap full” reading that did not reflect a real memory-pressure condition, clearing only after the operator manually cycled swap — BOTH guards refused on a PROXY signal that was NOT the real condition. Every GUARD / GATE / PRE-FLIGHT (host-safety, resource-budget, readiness, lock, capability, availability) MUST assert the REAL condition it CLAIMS, resolved from the AUTHORITATIVE source (the real `/proc/<pid>/cmdline`, the real cgroup/meminfo pressure metric, the real lock-holder liveness via `kill -0`, the real device identity), NEVER from a proxy signal that something which is NOT the condition can satisfy: (1) **a FALSE-POSITIVE REFUSAL is a FAIL-bluff (§11.4.1)** — blocking work when the condition is ABSENT is exactly as forbidden as passing while a defect is present (it halts real work AND teaches operators/agents to bypass guards); (2) **a FALSE-NEGATIVE PASS is a PASS-bluff (§11.4);** (3) **SELF-VALIDATED (§11.4.107(10))** — every

guard ships a golden-TRUE fixture (condition present → guard FIRES) AND a golden-FALSE fixture INCLUDING the known CARRIER/decoy case (condition absent but proxy signal present → guard MUST NOT fire), both wired into the meta-test; a guard firing on its golden-FALSE fixture is ITSELF the defect; (4)

conservative-safe default on an unresolvable signal (§11.4.101/§12) — refuse the destructive/high-blast-radius action and say so HONESTLY (“could not resolve the condition” is NOT a false positive; claiming a condition never verified IS one, §11.4.6); (5) **every refusal reports its RESOLVED evidence** (real pid+cmdline, real metric+threshold, real lock holder) so a false positive is diagnosable in one step. GENERALISES §11.4.196(D) (resource-detection-by-real-identity — to EVERY guard) + §11.4.180 (liveness PROVEN before reaping — to every guard’s condition). Classification: universal (§11.4.17). Composes §11.4.1/.6/.101/.107(10)/.109/.111/.180/.196(D)/§12/§12.6/§12.8/§12.12.

Propagation gate `CM-COVENANT-114-201-PROPAGATION` (literal `11.4.201`) + recommended gate `CM-GUARD-ASSERTS-REAL-CONDITION` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.201. Non-compliance is a release blocker. No escape hatch — no `--proxy-signal-OK`, `--substring-match-guard`, `--false-positive-refusal-acceptable`, `--unvalidated-guard`, `--refuse-without-evidence` flag.

EXTENSIONS landed 2026-07-15 (existing anchors strengthened — full text in `Constitution.md`). §11.4.89(G) subagent long-op discipline + conductor-owns-inherently-long-ops: a SUBAGENT MUST NOT run a single synchronous op exceeding the runtime’s no-progress stream watchdog (~5 min) — it backgrounds + polls (each poll IS the progress the watchdog observes); inherently-long ops (full builds / full suites / device flashes) are the CONDUCTOR’s to run (no such watchdog); a stall-killed subagent is a CRASH per §11.4.147 (respawned, never absorbed as a completion); watchdog budget read from the runtime, never guessed (§11.4.6). **§11.4.106(F)** write-seam hook enforcement: the doc/DB sync MUST be enforced at the COMMIT seam (hook refuses a commit whose staged set touches a chain source while its exports/DB are stale) + the BUILD seam (pre-build gate) + the CONSTITUTION-PULL seam (§11.4.164 post-update hook) — a stale CONTINUATION (§12.10) / session-resumption file (§11.4.131) / memory artefact while live state moved is a violation; honest boundary §11.4.6 — seam hooks give eventual-consistency-at-every-write, NOT literal real-time (that needs a watch daemon, carried as a tracked §11.4.197 upgrade path, never claimed as shipped). **§11.4.147(e)** API-QUOTA / RATE-LIMIT EXHAUSTION is a FIRST-CLASS CRASH

CLASS: an agent killed mid-work by a weekly/session/subscription quota boundary is `crashed` (never `complete`) — record the REAL limit signal + reason-class per §11.4.196(B), REBIND to the next OPERATIONAL alias per §11.4.196(A)/(C) and RESPAWN there (bounded-PARK per §11.4.101 only when EVERY alias is cooled), RESUME from the killed agent's captured transcript + partial state at its EXACT last point, and the §11.4.87/.94/.97/.126 loop done-condition MUST NOT read satisfied while any quota-killed entry is non-`complete` . **§11.4.196(F) CONFIGURED ≠ IN USE: forensic FACT** — the alias orchestrator was installed with a valid config yet its binding registry was EMPTY while THREE live workers ran outside it, so `promote` rebound ZERO tracks; EVERY spawn path MUST bind THROUGH the orchestrator, and the mechanism's readiness is its RUNTIME SIGNATURE (§11.4.108) — **the binding-registry rows MUST EQUAL the live-worker set** — an empty/partial registry alongside live workers is a FALSE-READY state and claiming "ready/in use" on config alone is a §11.4/§11.4.6 orchestration-layer bluff.

§11.4.142 + §11.4.194 re-affirmed (2026-07-15 audit, no change required): the operator re-insisted 3× this session that EVERY change passes an exhaustive code review, ALWAYS, before build/commit. Audited: §11.4.142 already binds EVERY change to ANY governed repository — "no exception ... no 'just a doc edit' exemption, no 'just a one-liner' exemption, no §11.4.92-self-review substitution, no 'trivial change' carve-out ... BEFORE it is accepted, committed, or built" — with independence load-bearing (§11.4.70/§11.4.20); §11.4.194 supplies the exhaustive all-scenario / prove-every-assumption / cross-check-captured-runtime-evidence precision bar; §11.4.134 iterates the review to a ZERO-finding, ZERO-warning GO; §11.4.125 places the gate before the pre-build sweep + main build. NO carve-out exists in any of them — the family is complete as written and needed no widening this round (§11.4.6: audited, not assumed).

§11.4.202 — Reporting directives: a report MUST auto-create a fully-populated, fully-synced workable item — never a prose acknowledgement (User mandate, 2026-07-15). Verbatim operator mandate: introduce reporting directives `ISSUE` / `BUG` / `TASK` that create a proper workable item filled with all details, with the WHOLE workable-items system automatically up to date and fully in sync — the main DB, all documents, and all external tracking systems; universal, reusable on every project, fully decoupled, recognized immediately by every project that pulls the constitution. Every governed project MUST expose REPORTING DIRECTIVES — the §11.4.140 registered actions `BUG` (Type=Bug), `TASK` (Type=Task), and `ISSUE` (generic report, CLASSIFIED into the §11.4.16 closed

set) — such that a plain-language report from the operator is TURNED INTO a real, tracked, fully-synced workable item, automatically, every time. A report that is discussed, acknowledged, answered in prose, or “noted” WITHOUT landing a tracked item is a §11.4.202 violation of §11.4 PASS-bluff severity at the requirements-intake layer: the requirement was accepted and then silently evaporated (the §11.4.197 loss-of-requirements failure, at its entry point). **(1) GRAMMAR (§11.4.140).** The three directives are ordinary registry rows, so EVERY §11.4.140 form works out of the box (`BUG :: <text>` , `DEFAULT::BUG :: <text>` , `/DEFAULT::BUG <text>` , `BUG ---> <text>`), PLUS a SIXTH form this anchor adds — the **single-colon form** `NAME: <text>` (`ISSUE: subtitles are late`), the shape operators actually type. The single-colon form is **REGISTERED-ACTION-ONLY by construction**: a single-colon token that is NOT a registered action is a NO-OP (an ordinary prompt), NEVER an ASK and NEVER a sub-system shortcut — because ordinary English prose (`TODO:` , `WARNING:` , `FIXME:`) shares that shape (`NOTE:` is now a registered §11.4.140 severity marker), and routing it to the §11.4.66 clarify path would question every such sentence (§11.4.6 honest boundary: the other five forms keep their ASK-on-unknown behaviour — their shapes do not occur in prose). Host-command collisions are handled by the EXISTING §11.4.140 conflict mechanism (`slash_bare: auto + slash_conflicts: [...]`), which already satisfies the operator’s explicit-prefix requirement: a bare `/NAME` that collides with a built-in (e.g. the documented Claude Code built-in `/bug`) is NOT honored, while the namespaced `/DEFAULT::NAME` form is ALWAYS unambiguous — and the `NAME: / NAME :: / NAME --->` forms are never affected by any host command.

(2) THE ITEM. The created item MUST carry, on ALL surfaces: a valid Status (§11.4.15, `Queued` at intake) + a Type from the CLOSED set {Bug | Feature | Task} (§11.4.16 — `ISSUE` is a REPORTING CHANNEL, never a fourth type; inventing one is a violation) + a stable auto-incremented id (§11.4.54), and a COMPREHENSIVE structured description (§11.4.148 D2 / §11.4.171): what / affected-scope / reproduction / acceptance. An undetermined section is recorded as an explicit `UNKNOWN:` gap — NEVER invented (§11.4.6). For `ISSUE` , the type is classified from the report’s content and stated as FACT; if the content does not determine it, the operator is ASKED (§11.4.66 / §11.4.105) BEFORE creation; ONLY when running autonomously where asking is impossible (§11.4.101) may the §11.4.16 lowest-stakes ambiguity default `Task` be used — and then the defaulted classification MUST be recorded verbatim in the item and surfaced for

reclassification, never silently asserted. **(3) THE FULL SYNC (the load-bearing half).** Creating the item is NOT sufficient: the directive MUST drive the WHOLE workable-items system into sync, in one shot — (a) the item lands in the SQLite single-source-of-truth (§11.4.93 / §11.4.95); (b) EVERY derived document is regenerated FROM the DB — the open + closed trackers, their summaries, and their `.md` / `.html` / `.pdf` / `.docx` siblings (§11.4.12 / §11.4.53 / §11.4.65 / §11.4.106); (c) the item is pushed to EVERY configured external tracker (§11.4.148 D5). **(4) NEVER A FAKED PUSH (§11.4.10 / §11.4.6 / §11.4.3).** A tracker whose credentials or whose client are absent is SKIPPED with an honest machine-readable reason (`credentials_absent` — names of the unset variables ONLY, never a value; `tracker_client_absent` — PENDING-OPERATOR-INPUT). A tracker PASS may be recorded ONLY after a real push command really exited 0. Fabricating, assuming, or silently omitting a push is a §11.4 bluff at the integration layer; a missing tracker NEVER blocks the item from landing in the DB + docs (the report is never lost because a mirror is unreachable). **(5) DECOUPLED ENGINE, CONSUMER-OWNED DATA (§11.4.28 / §11.4.177).** The item-creation + full-sync engine lives in the constitution submodule (`constitution/scripts/reporting/report_item.sh`), is inherited BY REFERENCE (never copied), and carries ZERO project literals; every project-specific value — DB path, id prefix, sync command, tracker commands + their required env vars — comes from a CONSUMER-OWNED config file (DATA, never an engine edit). The engine FAILS CLOSED with an actionable message when that config is absent — it never guesses a project's paths (§11.4.6). **(6) ANTI-BLUFF (§11.4 / §11.4.107(10)).** The mechanism ships four-layer coverage: a pre-build gate (`CM-REPORTING-DIRECTIVES`) whose FUNCTIONAL invariant SOURCES and RUNS the parser (never a grep-only assertion, §11.4.108), a paired §1.1 mutation test that proves each invariant is genuinely load-bearing, and an end-to-end suite run against a REAL SQLite DB with REAL binaries (no mocks, §11.4.27) that asserts a real row lands with the right Type/Status/id, that the docs are regenerated FROM the DB, and that an absent tracker SKIPS honestly — self-validated with a golden-good + golden-bad + negative-control fixture set (an oracle that passes its golden-bad fixture is itself the bluff). Classification: universal (§11.4.17) — the consuming project supplies its DB path, id prefix, sync command, and tracker bindings per §11.4.35. Composes §11.4.140 (grammar) / §11.4.93 / §11.4.95 (DB SSoT) / §11.4.15 / §11.4.16 / §11.4.54 (status + type + id) / §11.4.148 / §11.4.171 (comprehensive description) / §11.4.106 / §11.4.12 / §11.4.53 / §11.4.65 (doc sync) / §11.4.10 (credentials) / §11.4.66 / §11.4.105 (clarify, never guess) / §11.4.101 (autonomous decision) /

§11.4.197 (a started requirement never evaporates) / §11.4.28 / §11.4.177 (decoupling) / §11.4.164 (auto-registration on constitution pull) / §11.4.107(10) / §1.1. Propagation gate `CM-COVENANT-114-202-PROPAGATION` (literal `11.4.202`) + recommended gate `CM-REPORTING-DIRECTIVES` (registry declares ISSUE/BUG/TASK + the single-colon form is registered-only; the engine exists, is parse-clean, and is project-literal-free; the three directives really expand at runtime while prose does not; the honest-skip reasons exist and a tracker PASS is gated on a real exit 0) + paired §1.1 mutation (strip an action row, flip the single-colon form to non-registered-only, strip the parser branch, remove the prose NO-OP branch, make the engine fake a tracker push, or couple the engine to one project → the gate FAILs; strip the literal → the propagation gate FAILs).

Canonical authority: constitution submodule `Constitution.md` §11.4.202.

Non-compliance is a release blocker regardless of context. No escape hatch — no `--report-without-item`, `--prose-acknowledgement-OK`, `--skip-item-creation`, `--skip-tracker-sync`, `--fake-tracker-push`, `--invent-a-fourth-type`, `--hardcode-project-in-engine` flag.

§11.4.207 — Instant multi-stream resume engine: whole-fleet continuation state is a durable, content-addressed, atomically-committed snapshot resumed in $O(\text{changed})$ reads, not an $O(\text{history})$ re-read (User mandate, 2026-07-15). Compact summary: the continuation mechanism (§12.10 / §11.4.127 / §11.4.131 / §11.4.205) is EXTENDED — never replaced — by a mechanical engine so a fresh session resumes ALL work (every Track, main stream, and agent) at once, INSTANTLY and at token cost $O(\text{changed streams})$, by reading ONE snapshot manifest + one compact blob per stream, NEVER re-reading the full handoff history. (1) **Content-addressed Merkle store** — each stream's canonical state bytes are named by their sha256 (dedup: unchanged streams share a hash → a global fleet snapshot costs $O(\text{changed})$); a snapshot is ONE manifest referencing every stream's content-hash (§11.4.205(4)). (2) **$O(\text{streams})$ resume, $O(\text{history})$ NEVER on the resume path** — `RestoreAll` reads the manifest + one blob per stream; the append-only event ledger is audit + lost-update recovery only, never read to resume (§11.4.127). (3) **Tamper-evident + deterministic** — every blob is integrity-checked on read (bytes MUST hash to their id) and a deterministic round-trip (re-serialize must byte-match), no wall-clock in the hashed content, verified by a SELF-VALIDATING oracle: golden-good PASS + golden-bad detected FAIL + negative-control (a legitimately-older-but-valid snapshot) PASS — the false-

positive guard proving the verifier distinguishes a LAGGING copy from a TAMPERED one (§11.4.201/§11.4.107(10)/§11.4.206(3)). (4) **Durable + atomic + single-writer + crash-safe** — atomic ref write (temp → fsync → rename → dir-fsync, §11.4.205(6)), append-only ledger (§11.4.205(6)), single-writer-per-stream refusal (§11.4.206), advisory lock with provably-stale reap (`kill -0` liveness, never steals a live lock, §11.4.180/§9.2). (5) **Project-agnostic engine, consumer-owned data** — the engine (`github.com/vasic-digital/continuum` , GitLab mirror) carries ZERO project literals, fails closed rather than guess a store path (§11.4.6/§11.4.177), and is inherited BY REFERENCE as a depth-1 submodule under the §11.4.28(C) carve-out (`constitution/submodules/continuum/` , `helix-deps.yaml` , zero own-org deps); every consumer supplies its store path / actor / stream naming as DATA. Anti-bluff (§11.4/§11.4.108): the whole engine ships four-layer coverage (unit + integration + e2e + stress + chaos §11.4.85, race-clean) with a measured resume Metric (`blobs_read == 1 + streams` is the O(streams) proof) and the self-validating oracle as captured evidence (§11.4.5/§11.4.69/§11.4.107); honest boundary (§11.4.6): it snapshots the state an agent REPORTS, not its execution image (not a CRIU/durable-execution runtime, §11.4.112), and `est_tokens` is a documented `bytes/4` heuristic, not a tokenizer count. Classification: universal (§11.4.17). Composes §12.10 / §11.4.127 / §11.4.131 / §11.4.205 / §11.4.116 / §11.4.147 / §11.4.180 / §11.4.201 / §11.4.206 / §11.4.107(10) / §11.4.28 / §11.4.177 / §11.4.85 / §9.2. Propagation gate `CM-COVENANT-114-207-PROPAGATION` (literal `11.4.207`) + recommended gate `CM-CONTINUUM-RESUME-ENGINE-PRESENT` (the engine present under `constitution/submodules/continuum/` , `go test -race ./... green` , `continuum selfcheck` = good=PASS/bad=FAIL/negctrl=PASS, zero project literals) + paired §1.1 mutation (strip the literal → propagation gate FAILS; downgrade the resume path to read the ledger, or the oracle to pass its golden-bad → the mechanism gate FAILS; gate-code = separate work item). **Canonical authority:** constitution submodule `Constitution.md` §11.4.207. Non-compliance is a release blocker. No escape hatch — no `--resume-reads-history` , `--skip-integrity-check` , `--unvalidated-oracle-OK` , `--hand-typed-resume-state` , `--per-project-resume-reimpl` flag. **§11.4.208 — Operator-request-history document: every project maintains a project-local, always-in-sync ledger of every operator request/prompt (User mandate, 2026-07-15).** Compact summary: every governed project MUST maintain a single PROJECT-LOCAL, append-only, newest-first request-history document (e.g. `docs/requests/history.md` , path

declared once per §11.4.35) recording EVERY operator request/prompt with EXACTLY these operator-mandated fields — (1) full request content (verbatim where archived, else a faithful complete summary, marked which); (2) accepted timestamp with an EXPLICIT timezone stamp on every entry (project-declared default zone, `time UNKNOWN` where only a date is recoverable, never fabricated); (3) the Track that processed it (§11.4.176/.178/.182); (4) the live alias that took the item (§11.4.182/.196); (5) the model + effort used. (A) NO INVENTION (§11.4.6) — an unrecoverable field is recorded literally `UNKNOWN`, never guessed (a fabricated timestamp/alias/model is a §11.4/§11.4.1 requirements-record bluff); (B) honest reconstruction boundary for pre-mandate sessions (best-effort from durable dated records, unrecoverable fields `UNKNOWN`, boundary stated explicitly); (C) always in sync (§12.10) + §11.4.44 revision header + four-format export §11.4.65/§11.4.73; (D) keep-applying mechanism (a helper script and/or a `UserPromptSubmit` -class hook appends a newest-first row per new prompt with deterministic track/alias derivation §11.4.182, honest `? / UNKNOWN` when undeterminable; a helper-without-hook is honestly partial, hook-wiring a tracked §11.4.197/§12.10 follow-up, never claimed as automatic capture it does not perform); (E) complements NEVER replaces the per-session no-loss ledger (§11.4.197/.202), the workable-items DB SSoT (§11.4.93/.95), and the §12.10/§11.4.131 session-resumption artefacts (loss of a requirement at intake FORBIDDEN §11.4.197); (F) RULE-UNIVERSAL / DOCUMENT-PROJECT-LOCAL (§11.4.35) — this universal RULE lives in the constitution submodule, the DOCUMENT + its path + the default timezone are consumer-owned DATA and are NEVER placed into the constitution (§11.4.28). Classification: universal (§11.4.17). Composes §11.4.6/.35/.44/.65/.73/.118/.131/§12.10/.176/.178/.182/.196/.197/.202/.93/.95. Propagation gate `CM-COVENANT-114-208-PROPAGATION` (literal `11.4.208`) + recommended gate `CM-REQUEST-HISTORY-DOC-PRESENT` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.208. Non-compliance is a release blocker. No escape hatch — no `--request-history-optional`, `--skip-request-log`, `--no-timezone-stamp`, `--invent-missing-field`, `--request-history-in-constitution`, `--prompt-without-history-entry` flag.

§11.4.209 — Code-review MUST run on the Fable model at xhigh effort (Opus xhigh fallback) (User mandate, 2026-07-15). Compact summary: every mandatory independent code-review the constitution requires — the §11.4.125 code-review-agent gate before pre-build + main build, the §11.4.142 universal every-change-review, and the §11.4.194 exhaustive all-scenario/all-angle review,

iterated to a zero-finding GO per §11.4.134 — MUST be dispatched on the **Fable model at xhigh effort**; (A) fallback (closed, ordered) — ONLY when Fable is genuinely unavailable (FACT per §11.4.6, never guessed) use **Opus at xhigh effort**; a lower-effort or a non-Fable/non-Opus model is NOT an acceptable review substrate, xhigh REQUIRED in both cases; (B) independence preserved (§11.4.70/.20) — §11.4.209 sets the MODEL+EFFORT, not the independence / iterate-until-GO §11.4.134 / captured-evidence §11.4.5/.69/.107; (C) honest boundary §11.4.6 — a review claimed done that ran on a weaker model/lower effort than Fable-xhigh (or Opus-xhigh on genuine Fable-unavailability) is a §11.4 review-substrate bluff; the review's captured evidence records which model+effort actually performed it (and, on fallback, the Fable-unavailability fact). STRENGTHENS §11.4.125/.142/.194/.134, does not replace them. Classification: universal (§11.4.17). Composes §11.4.125/.134/.142/.194/.6/.20/.70/.196. Propagation gate CM-COVENANT-114-209-PROPAGATION (literal 11.4.209) + recommended gate CM-CODE-REVIEW-FABLE-XHIGH + paired §1.1 mutation. **Canonical authority:** constitution submodule Constitution.md §11.4.209. Non-compliance is a release blocker. No escape hatch — no --review-any-model, --review-lower-effort-OK, --skip-fable, --non-xhigh-review, --fallback-without-fable-unavailability-proof flag.

§11.4.208 — Operator-request-history document: every project maintains a project-local, always-in-sync ledger of every operator request/prompt (User mandate, 2026-07-15). Verbatim operator mandate: introduce a request-history document capturing per operator request/prompt — full content, accepted-timestamp with explicit timezone, the track that processed it, the alias that took the workable item, and the model + effort used — always in sync, no false data, no bluff; the operator's load-bearing correction (§11.4.35/§11.4.17): the RULE is universal (→ constitution submodule), the request-history DOCUMENT itself is PROJECT-LOCAL (never placed into the constitution). Every governed project MUST maintain a single project-local, append-only, newest-first **operator-request-history document** at a fixed project-declared path (§11.4.35, never moved without a §11.4.66 operator decision), each entry carrying EXACTLY: (1) **Request content** (verbatim where archived, else a faithful complete summary, marked which); (2) **Accepted (when)** — date + time + TIMEZONE stated EXPLICITLY on every entry (consuming project declares its default zone §11.4.35; date-only ⇒ time UNKNOWN, never fabricated); (3) **Track** (§11.4.176/.178/.182); (4) **Alias** (§11.4.182/.196); (5) **Model + effort**. (A) **NO INVENTION (§11.4.6)** — an

unrecoverable field is recorded literally `UNKNOWN` , never guessed; a fabricated timestamp/alias/model is a §11.4/§11.4.1 bluff at the requirements-record layer. (B) **HONEST RECONSTRUCTION BOUNDARY** — pre-mandate sessions are best-effort reconstructed from the durable dated record with unrecoverable fields marked `UNKNOWN` and the boundary stated explicitly, never silently implied complete. (C) **ALWAYS IN SYNC + EXPORTED** — §11.4.44 revision header, §12.10 append-on-every-new-prompt, §11.4.65/§11.4.73 four-format export (`.md + .html + .pdf + .docx` where applicable, siblings $\geq$ `.md`). (D) **KEEP-APPLYING MECHANISM** — a helper script and/or a `UserPromptSubmit` -class hook appends a newest-first row per NEW prompt at capture time (track/alias derived deterministically §11.4.182, honest `? / UNKNOWN`); a helper landed WITHOUT the auto-capture hook is honestly a partial mechanism (hook wiring = tracked §11.4.197/§12.10 follow-up, never claimed as auto-capture it does not perform). (E) **COMPLEMENTS, NEVER REPLACES** any per-session no-loss ledger (§11.4.197/.202), the workable-items DB SSoT (§11.4.93/.95), or the §12.10/§11.4.131 resumption artefacts; requirement loss at intake is FORBIDDEN (§11.4.197). (F) **RULE-UNIVERSAL / DOCUMENT-PROJECT-LOCAL (§11.4.35)** — the RULE lives in the constitution submodule; the document + its path + default timezone are consumer-owned project DATA (§11.4.28 decoupling), NEVER placed into the constitution. Classification: universal (§11.4.17). Composes §11.4.6/.35/.44/.65/.73/.118/.131/.176/.178/.182/.196/.197/.202/.93/.95/§12.10. Propagation gate `CM-COVENANT-114-208-PROPAGATION` (literal `11.4.208`) + recommended gate `CM-REQUEST-HISTORY-DOC-PRESENT` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.208. Non-compliance is a release blocker. No escape hatch — no `--request-history-optional` , `--skip-request-log` , `--no-timezone-stamp` , `--invent-missing-field` , `--request-history-in-constitution` , `--prompt-without-history-entry` flag.

§11.4.209 — Code-review MUST run on the Fable model at xhigh effort (Opus xhigh fallback) (User mandate, 2026-07-15). Verbatim operator mandate: the mandatory independent code-review MUST ALWAYS be performed with the **Fable model at xhigh effort**; if Fable is unavailable, use **Opus at xhigh effort**. Every mandatory independent code-review the constitution requires — the §11.4.125 code-review-agent gate before pre-build + main build, the §11.4.142 universal every-change-review (no exception), the §11.4.194 exhaustive all-scenario/all-angle review, iterated to a zero-finding GO per §11.4.134 — MUST be dispatched on the

Fable model at xhigh effort. (A) **Fallback (closed, ordered)** — ONLY when Fable is genuinely unavailable (established as FACT §11.4.6, never a guessed unavailability) the review runs on **Opus at xhigh effort**; a lower-effort or a non-Fable/non-Opus model is NOT an acceptable review substrate, **xhigh** is REQUIRED in both cases (the §11.4.194 all-angle depth is unreachable lower). (B) **Independence preserved (§11.4.70/.20)** — the Fable/Opus-xhigh model is the structurally-separated reviewer, never the author; §11.4.209 sets the review's MODEL + EFFORT, not its independence / iterate-until-GO (§11.4.134) / captured-evidence (§11.4.5/.69/.107). (C) **Honest boundary (§11.4.6)** — a review claimed done that ran on a weaker model / lower effort than Fable-xhigh (or Opus-xhigh on genuine Fable-unavailability) is a §11.4 bluff at the review-substrate layer; the review's captured evidence records which model + effort actually performed it (+ the Fable-unavailability FACT on fallback). STRENGTHENS §11.4.125/.142/.194/.134 — does NOT replace them; it pins the review's execution substrate. Classification: universal (§11.4.17). Composes §11.4.125/.134/.142/.194/.6/.20/.70/.196. Propagation gate `CM-COVENANT-114-209-PROPAGATION` (literal `11.4.209`) + recommended gate `CM-CODE-REVIEW-FABLE-XHIGH` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.209. Non-compliance is a release blocker. No escape hatch — no `--review-any-model`, `--review-lower-effort-OK`, `--skip-fable`, `--non-xhigh-review`, `--fallback-without-fable-unavailability-proof` flag.

§11.4.210 — Zero-loss request/prompt intake: every operator request/prompt MUST be mechanically CAPTURED, TRACKED, and PROCESSED — never skipped, ignored, avoided, or lost (User mandate, 2026-07-15). Verbatim intent: “every request/prompt we issue MUST ALWAYS be respected, taken into account, executed and processed! No avoiding, ignoring, skipping or any form of loss!” Every operator request/prompt MUST be (a) CAPTURED in the §11.4.208 request-history ledger MECHANICALLY at intake via a WIRED verified-runtime-state `UserPromptSubmit` auto-capture hook (NOT a manual step, NOT a “tracked follow-up” — §11.4.210 PROMOTES §11.4.208(D)'s hook from partial-follow-up to MANDATORY per the §11.4.205 enforced-not-advisory pattern; forensic FACT 2026-07-15: an entire release-drive session's prompts went un-recorded because the hook was left un-wired); (b) TRACKED → a §11.4.202 workable item / §11.4.197 completion, or an explicit evidence-backed no-op-with-reason; (c) PROCESSED — respected/executed, never dropped. Loss at ANY of the three stages is a §11.4 PASS-bluff at the requirements-intake layer of §11.4.197's

severity class. Docs-chain-bound (§11.4.106), gated (§11.4.75 — hook verified-installed in the live hook path + content-matched, ledger fresh); the autonomous loop (§11.4.126/.87/.94/.97/.103) processes every captured request (done-condition never satisfied while any captured request is un-processed, §11.4.147(e)); honest §11.4.6 (UNKNOWN never invented). Classification: universal (§11.4.17).

Composes §11.4.197/.202/.208/.205/.106/.126/.75/.6/.147(e). Propagation gate

CM-COVENANT-114-210-PROPAGATION (literal 11.4.210) + recommended gate

CM-REQUEST-CAPTURE-HOOK-WIRED + paired §1.1 mutation. **Canonical authority:**

constitution submodule Constitution.md §11.4.210. Non-compliance is a release blocker. No escape hatch — no --skip-request-capture , --manual-ledger-OK , --hook-optional , --drop-unprocessed-request , --ledger-may-be-stale flag.

§11.4.211 — Merge-conflict resolution during main → feature/product/flavor merges MUST run on the Fable model at xhigh effort (Opus xhigh fallback)

(User mandate, 2026-07-16). Verbatim operator mandate: “During the merge of main branch to all feature/product/flavor branches of tracks 2 - 4, if conflicts happen they MUST BE resolved using Fable model with xhigh effort! If Fable is not available, then with Opus with xhigh!” Any merge-CONFLICT resolution performed while integrating the canonical trunk (main / master) into a feature / product / flavor branch — the §11.4.188 regular main → feature merge cadence, the §11.4.195 branch-taxonomy integration, the §11.4.41 merge-first pipeline, the §11.4.167(D) trunk-into-stream sync AND the §11.4.167(I) approval-gated back-merge — MUST be carried out on the **Fable model at xhigh effort**; the merge-resolution analogue of §11.4.209 (which pins the same model + effort for code-REVIEW), pinning it for the DECISION-MAKING of conflict resolution (which side to keep, how to UNION both sides without loss, marker removal, semantic reconciliation of divergent edits). **(A) Fallback (closed, ordered):** ONLY when Fable is genuinely unavailable (not configured / rate-limited / unreachable, established as FACT §11.4.6 — never a guessed unavailability) the conflict resolution runs on **Opus at xhigh effort**; a lower-effort or a non-Fable/non-Opus model is NOT an acceptable substrate for resolving merge conflicts, xhigh REQUIRED in both cases (a careless resolution silently drops a commit or commits a marker). **(B) Scope:** the operator named tracks 2–4 (feature / product / flavor); GENERALISES per §11.4.17 to EVERY controlled main → (feature|product|flavor) merge on EVERY track that hits a conflict. **(C) Safety preserved:** sets the MODEL + EFFORT of the conflict-resolution work only, BOUNDED BY

§11.4.113 (NEVER force-push) / §9 / §9.2 (no commit lost, hardlinked pre-op backup before any large/risky merge) / §11.4.41 step 3 (ZERO conflict markers committed, ZERO file silently dropped, union of both sides preserved) / §11.4.188(4) anti-bluff (post-merge smoke GREEN + marker-scan empty + no-lost-commit) / §11.4.84 (quiescent-only) / §11.4.37/§11.4.71 (fetch-first). **(D) Honest boundary (§11.4.6):** a merge whose conflicts were resolved on a weaker model / lower effort than Fable-xhigh (or Opus-xhigh on genuine Fable-unavailability) is a §11.4 bluff at the merge-resolution-substrate layer — the merge’s captured evidence records which model + effort resolved the conflicts (and, on fallback, the FACT of Fable’s unavailability). STRICT EXTENSION of §11.4.209 to merge-CONFLICT resolution; STRENGTHENS §11.4.188/§11.4.195/§11.4.41 by pinning the substrate of their conflict-resolution step. Classification: universal (§11.4.17). Composes §11.4.209/.188/.195/.41/.113/§9/§9.2/.6/.167(D)(I)/.84/.37/.71/.196. Propagation gate `CM-COVENANT-114-211-PROPAGATION` (literal `11.4.211`) + recommended gate `CM-MERGE-CONFLICT-FABLE-XHIGH` + paired §1.1 mutation. **Canonical authority:** constitution submodule `Constitution.md` §11.4.211. Non-compliance is a release blocker. No escape hatch — no `--merge-conflict-any-model`, `--skip-fable-merge-resolution`, `--non-xhigh-merge-conflict`, `--resolve-conflict-lower-effort-OK`, `--fallback-without-fable-unavailability-proof` flag.

§11.4.212 — Main README is the canonical starting-point / entry point for ALL project documentation — no doc may be an orphan unreachable from README (User mandate, 2026-07-16). Verbatim operator mandate: “Manual tests catalog must have linking all through the main README as the starting point of this and any other Project documentation! This fact and mandatory rule MUST be added into the constitution.” The main README (`README.md` at the repository root) is the canonical STARTING POINT / entry point for ALL project documentation, for THIS project and EVERY project incorporating this constitution. Every project document in the §11.4.65 export scope — INCLUDING the manual-tests catalog (the §11.4.185 manual-QA testing catalog), the §11.4.153 per-feature Status / Status_Summary docs, the Issues / Fixed trackers + their summaries (§11.4.12 / §11.4.53 / §11.4.57), and every guide / research note / changelog / script companion doc — MUST be reachable / linked (DIRECTLY or TRANSITIVELY) from the main README. No project document may be an ORPHAN — a doc that no link-path from README reaches is a §11.4.212 violation. STRICT GENERALISATION of §11.4.57 (which mandated a `Tracked-`

Items + Status Documents doc-link section covering the tracker / status set) to EVERY §11.4.65-scope document, so the README is the single navigable root of the project's documentation tree. **(A) Reachability = a link path.** README → guide → sub-doc is fine (transitive); the requirement is that SOME path from README reaches every in-scope doc, not that every doc is linked directly at top level. **(B) Always-sync (§11.4.59 / §11.4.106 docs-chain).** A newly-added §11.4.65-scope doc that no README-reachable path links is a §11.4.212 violation the moment it lands; the manual-tests catalog specifically MUST be linked from the README (the operator's load-bearing example). **(C) Honest boundary (§11.4.6).** Reachability is a link-graph property (every doc FINDABLE from README), NOT a claim about each doc's CORRECTNESS or freshness (those stay §11.4.44 / §11.4.106 / the doc's own gates). **(D) Composition with §11.4.57.** §11.4.57's Tracked-Items section remains the mandated README landing block for the tracker / status set; §11.4.212 adds that EVERY OTHER in-scope doc is likewise reachable — §11.4.57 is the specialisation, §11.4.212 the generalisation. Classification: universal (§11.4.17) — the consuming project supplies its README path, its manual-tests-catalog path, and its §11.4.65 doc-set enumeration per §11.4.35. Composes §11.4.57 / §11.4.59 / §11.4.65 / §11.4.106 / §11.4.153 / §11.4.185 / §11.4.44 / §11.4.12 / §11.4.53. Propagation gate CM-COVENANT-114-212-PROPAGATION (literal 11.4.212) + recommended gate CM-README-DOC-ENTRYPOINT-COMPLETE + paired §1.1 mutation. **Canonical authority:** constitution submodule Constitution.md §11.4.212. Non-compliance is a release blocker. No escape hatch — no --orphan-doc-OK, --readme-not-entrypoint, --skip-readme-link, --doc-unreachable-OK flag.

§11.4.213 — FEATURE research-scheduling directive: recognized via all §11.4.140 forms, SCHEDULES (never synchronously executes) a deep, enterprise-grade research + implementation-planning effort as a tracked workable item (operator-directed mandate, 2026-07-16). Compact summary: extends the §11.4.202 reporting-directive family with a fourth registered action FEATURE — recognized via ALL SIX §11.4.140 grammar forms (FEATURE :: x / DEFAULT::FEATURE :: x / /FEATURE x / /DEFAULT::FEATURE x / FEATURE ---> x / the single-colon FEATURE: x , registered-action-only so lowercase prose like "feature: ..." never expands and is never asked about, §11.4.6) — that SCHEDULES rather than synchronously executes a deep research + implementation-planning effort: the directive creates a Type=Task / Status=Queued workable item (§11.4.15/16/.54) by INVOKING the SAME §11.4.202

`report_item.sh` engine (never a duplicate implementation of the item-creation/
 DB-sync/tracker-push machinery — a second divergent tracker-push path would
 itself be a bluff-surface regression) and appends it to a durable `docs/requests/
 feature_queue.md` queue (mirroring `BACKGROUND` 's `docs/requests/
 background_queue.md` pattern) so a scheduled request is never dropped; the
 item's comprehensive description (§11.4.148/.171) embeds the FULL research
 mandate as its acceptance-defining work program — deep web research on best
 incorporation (§11.4.8/.99/.150), systematic-debug of every enumerated weak-spot/
 gap/danger-zone with a risk-free rock-solid design per gap (§11.4.102), always
 enterprise-grade/bleeding-edge/innovative, exhaustive documentation down to
 lines-of-code + micro-POCs + diagrams + SQL + templates (§11.4.65/.73), reuse-
 first investigation of vasic-digital/HelixDevelopment components kept decoupled
 (§11.4.28/.74/.177), full test-type + Challenges + HelixQA planning (§11.4.27/.169),
 fine-grained phased/task/subtask breakdown, enterprise-scalability + max-
 performance planning, OpenDesign UI/UX wireframes where applicable
 (§11.4.162/.190), full CodeGraph integration (§11.4.78/.79/.80), and fully-synced
 workable items across the SQLite SSOT + every connected external tracker
 (§11.4.93/.95/.148 D5), honestly SKIPPING any absent tracker (§11.4.10) —
 NEVER a faked push. The multi-day research itself is EXECUTED by the standing
 autonomous loop (§11.4.87/.94/.97/.103/.126) when it claims the item, driven to a
 genuinely COMPLETED-and-wired or explicitly evidence-backed CLOSED terminal
 state under §11.4.197 — never left un-wired in the backlog; “scheduled” is never
 reported as “done” (§11.4.6). Anti-bluff four-layer coverage: pre-build gate `CM-
 FEATURE-DIRECTIVE` (FUNCTIONAL — sources + RUNS the §11.4.140 parser,
 never grep-only, §11.4.108) + paired §1.1 mutation + a real-DB end-to-end
 evidence trail (temp SQLite DB, no mocks, §11.4.27). Classification: universal
 (§11.4.17). Composes
 §11.4.140/.202/.197/.8/.99/.150/.102/.65/.73/.28/.74/.177/.27/.169/.162/.190/.78/.79/.80/.93/.95/.14
 Propagation gate `CM-COVENANT-114-213-PROPAGATION` (literal 11.4.213) +
 recommended gate `CM-FEATURE-DIRECTIVE` + paired §1.1 mutation. **Canonical
 authority:** constitution submodule `Constitution.md` §11.4.213. Non-compliance
 is a release blocker. No escape hatch — no `--feature-runs-synchronously` ,
`--skip-feature-queue` , `--feature-without-research-mandate` , `--
 reimplement-tracker-push` , `--hardcode-project-in-engine` , `--skip-item-
 creation` flag.

EXTENSIONS landed 2026-07-17 (existing anchors strengthened — full text in `Constitution.md` ; research-derived from a consuming project's status-without-custody forensics). Forensic FACTs (genericised, all measured 2026-07-17): 576 tracked items, 182 claiming a done/ready status, **173 (95%) with NO registered regression guard** — an operator manually sampled 6 done-claiming items and found **6/6 broken**, the expected draw of that arithmetic, not bad luck; recorded reopen-per-recorded-fix **52%** (published elite <10%), itself an UNDERCOUNT because recurrences re-entered the tracker as NEW ids and statuses moved with ZERO history rows (one item: 4 Reopened events, 0 terminal events); the standing guard suite's exit code read ONLY its FAIL count — measured live on a release-candidate target: **61 registered guards** → **PASS 0 / FAIL 5 / PENDING 56** → **exit 0 “not blocked”** — and the release-tag tooling consulted the guard suite **ZERO times**, so to the release gate “never verified” and “verified good” were the same colour; the project's own history selected the discriminator — every fix whose done-claim was preceded by an on-device RED → GREEN polarity flip that HELD recorded zero reopens, and every fix confirmed at source-green bounced. **No new anchor number was minted**: §11.4.115 and §11.4.146 already said what the operator mandates and did not bind because they are prose consulted by agents, not conditions checked by seams — the deliverable is the MECHANICAL BINDING (the §11.4.205 test applied: “if I documented this and nobody implemented the hook, would anything catch it?”). **§11.4.115(F) — machine-written verdict pairs + guard-viability**: RED and GREEN are harness-written verdict files, never prose — carrying the item id, guard identity, polarity, exit code, the target's artifact fingerprint READ FROM THE TARGET AT RUN TIME, iterations (≥ 3 , §11.4.50), and evidence files whose CLASS matches the defect's layer (a source-grep transcript can never satisfy a user-visible/pixel/audio class — the wrong-layer oracle dies by construction); GREEN requires a prior RED with exit $\neq 0$ on the PRE-fix artifact and a DIFFERENT fingerprint (identical fingerprints prove the fix was never deployed); a guard never observed FAILing on the genuinely-broken artifact is unvalidated instrumentation and mints no verdicts (a guard structurally unable to FAIL — or to PASS — satisfies nothing); the canonical §1.1 mutation for a landed fix is the fix-commit's REVERT, and a mutation whose diff only deletes the literal strings the gate greps for is a refused tautology. Recommended gate `CM-GUARD-VERDICT-MACHINE-WRITTEN` . **§11.4.135 (verdict-coverage extension) — ABSENCE of a verdict blocks exactly as a FAIL does**: the suite's exit semantics MUST incorporate COVERAGE, never only the FAIL count; the release seam refuses the tag while `uncovered = registered ^`

topology-present $\wedge$ no-verdict-for-the-candidate-fingerprint is non-empty OR any FAIL exists; the release-tag tooling MUST consult the suite/verdict store mechanically; REQUIRES-REBUILD PENDING is honest on intermediate artifacts and self-clearing on the candidate (re-run on the candidate, or the item leaves the release's done-set — both forced, neither a permanent exemption); topology-absent guards are exempt ONLY via a checked-in target-pool topology map and MUST be enumerated in the release changelog as honest gaps (§11.4.3/§11.4.118) — never silent; a naive “PENDING always blocks” is REJECTED as a §11.4.201 FAIL-bluff; pre-existing unguarded backlogs adopt via a ONE-TIME monotone-decrease ratchet snapshot so day one is green and the disease cannot spread. Recommended gate CM-RELEASE-VERDICT-COVERAGE (self-tested with golden-good / golden-bad / negative-control per §11.4.205/§11.4.107(10)).

§11.4.146(D3) — status-custody mechanical binding: a done/ready status write is REFUSED without the file chain registry row keyed by the item's EXACT stable id (checked-in alias map for legacy keys — a guard keyed to a non-item id fails the join and satisfies nothing) → executable registered guard → §11.4.115(F) RED+GREEN verdict pair → class-matched evidence — prose appears nowhere in the chain; Seam A = the workable-items engine (§11.4.93) refuses the write (project-agnostic per §11.4.28, custody paths are consumer DATA per §11.4.35); Seam B = a FULL-TABLE sweep gate on every pre-build run (not a diff check — catches raw-DB writes that bypass Seam A); Seam C = the §11.4.135 release seam; §11.4.205 construction order — the executable sweep + its golden-bad meta-fixture (a temp SSoT overlay with a chain-less terminal item; the meta-test RUNS the sweep and requires FAIL, and asserts the sweep is WIRED) land in the SAME commit as any governance text citing them. Recommended gate CM-STATUS-CUSTODY . Honest boundary (§11.4.6), carried explicitly: the binding proves no status outruns its evidence and the REPORTED defect cannot silently return under the guard's conditions — it does NOT prove the feature bug-free (family dedup §11.4.186 remains), does NOT make a miscalibrated oracle honest (§11.4.107(10) goldens are necessary, not sufficient), does NOT replace the §11.4.185 manual-QA sufficiency gate, and makes target availability a HARD dependency of status progress BY DESIGN (a visible block replaces a silent pass). All Classification: universal (§11.4.17) — no hardware/vendor/project literal; consumers supply their registry, verdict-store, sweep, ratchet, and topology-map

paths as DATA per §11.4.35. Propagation gates for the literals 11.4.115 / 11.4.135 / 11.4.146 are unchanged and stay GREEN; the three recommended mechanism gates + paired §1.1 mutations are gate-code = separate work item.

EXTENSIONS + NEW ANCHORS landed 2026-07-17 (round 2 — the carrier/measurement-integrity harvest). Research-derived from one cycle's forensics in a consuming project, where EVERY false finding had ONE shape: **a token that MENTIONS X was matched as X, or an instrument returned silence that was read as data.** Genericised FACTs: a removal's completeness was measured by counting references to the removed component and the count matched the AUDIT-TRAIL COMMENTS documenting the removal (a completed removal read as failed); a pattern's escapes were eaten by an intervening shell so the far side searched a LITERAL and matched nothing (a false "the component is missing" that would have implied the target could not boot — it had booted); `\|` under an extended-regex dialect is a LITERAL pipe so a mandated-fields check searched a string nobody wrote and nearly reported the fields absent (all were present); a `check | head` pipeline returned the PAGER's exit status so a parse-check reported "clean ✓" three times without checking anything; a line-start-anchored pattern missed an indented, colour-escaped line and read a present verdict as no-verdict; a binary extractor split a multi-byte character and reported a present string as absent. §11.4.196(D)'s process-carrier and §12.12's self-match had ALREADY anchored this shape TWICE and it kept recurring — that recurrence is the argument for the general rule. The self-demonstrating instance: a mutation-residue scan run over the governance repo flagged the very anchor that DEFINES the mutation markers, because that anchor quotes them in its rule text (proven a carrier, not residue, by an identical HEAD-vs-worktree count). **The fix, proven in-session: the CONTROL NEEDLE** — an artifact check returned 0 hits; instead of reporting absence it ran a string KNOWN present through the SAME path, got 1 hit, and only THEN reported the zero as real.

§11.4.201 (6)(7)(8) — scope widened to EVERY load-bearing measurement + carrier-vs-thing + control needle + metric-validation. (6) clauses (1)-(5) bind guards (decision seams that refuse/allow) with a refusal-shaped taxonomy; (6)-(8) bind **every measurement whose result is load-bearing** (grep/count/parse/exit-code/scan/metric/artifact-inspection) with a MEASUREMENT-shaped taxonomy — a **FALSE-NULL** (a zero read as absence when the instrument could not see) and a **FALSE-MATCH** (a hit on a carrier) are §11.4/§11.4.1 bluffs at the measurement layer; the null is more dangerous because a broken instrument and a clean artifact

return the identical quiet zero. (7)(a) **MATCH STRUCTURE, NOT SUBSTRING** — distinguish the THING from a CARRIER that MENTIONS it. (7)(b) **A NULL IS NOT EVIDENCE until a CONTROL NEEDLE proves the instrument can see** — same instrument, same path, same artifact, a known-present needle **MUST** return non-null; needle absent $\Rightarrow$ the instrument is blind and the zero says NOTHING (report the blindness, never the absence §11.4.6). **The needle MUST share the certified query's LOAD-BEARING FEATURES** (same dialect constructs / anchoring / quoting / encoding): a bare-literal needle certifies only the layers a bare literal crosses — the (7)(c) dialect **FACT** refutes the naive form directly, since a literal needle sails through the same tool+path and would 'certify' a zero from a query still searching a string nobody wrote, whereas a needle carrying the **SAME** alternation dies with it. Earned inference is therefore **BOUNDED**: needle found $\Rightarrow$ the path can see the needle's **QUERY CLASS** $\Rightarrow$ a zero from that same class is evidence; a zero from a query using features the needle never exercised needs a needle of its own class. Generalises §11.4.115(F)'s validate-the-detector from guards to every measurement, **PER-QUERY** — **golden fixtures cannot substitute** (§11.4.146(D3): goldens necessary-not-sufficient), an instrument can pass every fixture on the authoring host and be blind on the real path. (7)(c) **THE PATH IS PART OF THE INSTRUMENT** — the tool **PLUS** every layer the query crosses (wrapper-shell quoting, regex dialect, pipeline exit status, anchoring, encoding); each returns a **CLEAN CONFIDENT WRONG** answer without crashing, so §11.4.1's script-bug clause (which covers crashes-into-FAIL) never fires and §11.4.67's parse-check (parse-time only) cannot see it — the control needle is the **ONLY** mechanism that tests the path. (8) **A METRIC MUST BE VALIDATED AGAINST THE DEFINITION OF DONE** — construct the correct end-state and evaluate the metric there; if the correct end-state does not move it to target, the metric is **INVALID** and **MUST NOT** gate work (a metric the correct outcome **REFUTES** will, if enforced, drive work **AWAY** from correctness); honest boundary §11.4.6 — this constitution already ships one such gating metric (§11.4.135's adoption ratchet, blocking at the **RELEASE-TAG** seam) whose validation against the definition of done is **OWED**, not claimed, and which every consuming project operating a ratchet **MUST TRACK** as a §11.4.197 item (stated as the obligation it is — this anchor asserts **NO** project has already filed it; asserting an unverified tracker state would be the §11.4.6 violation the clause forbids). **Interim precedence:** until that validation lands the §11.4.135(5) ratchet **KEEPS GATING** — clause (8) does **NOT** disarm it (an unvalidated ratchet blocking a release is bounded + visible + operator-resolvable; disarming the only adoption mechanism would silently re-open

the unguarded-backlog hole §11.4.135 closes — the §11.4.101 reversible-safe choice); clause (8) binds every NEW metric from this anchor forward, and the ratchet's validation as owed work — never as a licence to switch it off. Gate contract landed in the canonical file: CM-GUARD-ASSERTS-REAL-CONDITION EXTENDED to the measurement layer ((i) a reported absence cites a class-matched control-needle result; (ii) the needle's class matches the certified query's load-bearing features) + NEW recommended sibling CM-METRIC-VALIDATED-AGAINST-DONE (every gating/ranking metric carries a recorded definition-of-done evaluation reaching target) + paired §1.1 mutations (bare-literal needle on a dialect-dependent query → (ii) FAILs; needle stripped → (i) FAILs; metric whose correct end-state moves AWAY from target → FAILs). Gate-code = separate work item — the contract is landed, the code is NOT claimed shipped (§11.4.6).

§11.4.112 (5) — impossibility verdicts are BOUNDED and MUST say where the edge is. FACT: a TRUE verdict (“relocating another application’s protected video to a secondary display *while that application’s UI stays in place, without its cooperation* is impossible”) LEAKED to an ADJACENT VIABLE goal (whole-task placement) for six weeks, while the platform had shipped a public permission-free SDK call for the adjacent goal several major versions earlier and its own connected-display docs PRESCRIBED that call. The anchor was ASYMMETRIC: clause (3) hardens a verdict against reopening while nothing bounded its REACH. Now every structurally-impossible verdict MUST (a) STATE ITS EXACT SCOPE (the narrowest claim the evidence supports — every load-bearing qualifier is part of the verdict; a qualifier dropped is a verdict widened), (b) ENUMERATE THE ADJACENT GOALS IT DOES NOT COVER (unexamined neighbour ⇒ UNKNOWN + tracked follow-up, never silently absorbed), (c) BE CITATION-FENCED (citing it beyond its fence, in planning/investigation/design/status, without its OWN §11.4.150 research + proof is a §11.4.6 violation + a verdict-layer bluff of the class §11.4.199 names). Honest boundary — outside the fence ≠ viable, it means UNDECIDED, earned by its own evidence never by inheritance in either direction. Why a clause not a narrative: §11.4.150 already documented this class in prose and REQUIRES the research pass before a verdict is EARNED — the leak still happened, because §11.4.150 fires where a verdict is MINTED and this leak happened where one was CITED as a premise, crossing no closure seam. Gate contract landed: CM-WONT-FIX-STRUCTURAL-IMPOSSIBILITY EXTENDED (scope

statement + adjacent-goals-not-covered enumeration required; citing a verdict outside its stated scope without that goal's own §11.4.150 proof → FAIL) + paired §1.1 mutations. Gate-code = separate work item (§11.4.6).

§11.4.120 (seam-placement) — a gate whose precondition the gated work itself produces is on the WRONG SEAM. §11.4.120's trigger is a gate that FAILs (a gate that RAN); its sibling is a gate that CANNOT RUN in the imposed ordering, and the response is not reconciliation but re-seating. FACT: work was ordered “run checks E1-E4, then implement P1 only if they pass” — but the gate's own text defined E2 as running ON a P1-gated build, so E2 could not exist before P1; the source of truth gated ENABLING (building behind a default-off flag was permitted), and the gate had been imposed on AUTHORSHIP. (a) if gate G requires artifact A and A is produced by the work G gates, G belongs at a LATER seam (activation/enablement/release), never authorship — imposing it there is a deadlock, not rigor; (b) RE-READ the source of truth before re-seating (the gate's own text names its real seam and governs over any restatement — a restatement that moved it from enabling to writing is a §11.4.6 misstatement; obey the source, never weaken the gate); (c) the forbidden responses are §11.4.120's own (no fake-pass/delete/tautology, no build-then-retro-declare, and the work is NOT abandoned as “blocked” when the gate merely sits on the wrong seam). Honest boundary — an unrunnable gate is NOT automatically mis-seated (§11.4.102 first); it may be correctly seated with a genuinely-missing precondition = the §11.4.69 `artifact_not_yet_built` NOT-YET-RUNNABLE state. §11.4.110's clash detector cannot catch this class (its input is a DIFF; an ordering instruction is not a diff). The “report the circularity, don't improvise” half is ALREADY covered by §11.4.6/§11.4.66/§11.4.101/§11.4.201(4) — this clause claims only the seam-placement rule. Gate contract landed: `CM-GATE-RECONCILED-NOT-FAKE-PASSED` EXTENDED (an imposed gate whose precondition the gated work itself produces → FAIL naming the circularity, re-seat at activation) + a golden-FALSE fixture that MUST include a correctly-seated gate merely awaiting its artifact (§11.4.69 `artifact_not_yet_built`) which the gate MUST NOT fire on, else it becomes the §11.4.201(1) false positive + paired §1.1 mutation. Gate-code = separate work item (§11.4.6).

§11.4.69 (closed reason-set) — `artifact_not_yet_built` added. The NOT-YET-RUNNABLE class (check CORRECT, topology PRESENT, precondition artifact absent — fix not built/deployed) had NO legal code, so a §11.4.135 REQUIRES-REBUILD PENDING had to misuse a false code (`hardware_not_present` / `feature_disabled_by_config` — both untrue, a

§11.4.6 misstatement) or fail the helper at call time. Reporting it FAIL is a §11.4.1 FAIL-bluff (sends teams to fix healthy code); PASS is a §11.4 PASS-bluff. Its seam-dependence is §11.4.135's, NOT flat — honest+non-blocking on an intermediate artifact, BLOCKING on the release candidate; a flat “always blocks” is REJECTED per §11.4.135 as itself a §11.4.201 FAIL-bluff. **Audited, NO new anchor:** the operator-candidate “BLOCKED as a third verdict state” is ALREADY fully covered — §11.4.135's 2026-07-17 verdict-coverage extension states it more precisely (seam-dependent, and it explicitly rejects the flat framing), §11.4.3 forbids the PASS, §11.4.201(1) forbids the FAIL, §11.4.116/§11.4.45 already carry the verdict vocabulary; only the reason-code was missing.

§11.4.214 (NEW) — Recurrence-links-not-mints: a defect that returns MUST reopen its existing item, never enter as a new id. FACT: recurrences of already-“fixed” defects re-entered as BRAND-NEW ids (≥ 5 chains, one three deep), so the ORIGINAL's reopen counter never incremented — and the §11.4.55 `reopens_count`, the exact signal §11.4.132(d)/§11.4.189 use to rank the most-fragile work for the deepest scrutiny, **stays quiet precisely on the items that keep breaking**; the recorded 52% reopen rate was a KNOWN UNDERCOUNT for this reason; ~15 items mapped to ~9 real root causes, so “fix all 15” would have forked duplicate/conflicting fixes on one cause. Before minting a NEW id from ANY intake path, check for an existing item describing the SAME defect → RESOLVE the match THROUGH its `duplicate-of` links to the CANONICAL CHAIN HEAD (§11.4.90's closure vocabulary includes `duplicate-of`, and clause (5) MANUFACTURES such items — reopening one would resurrect a non-canonical copy beside its live canonical item, forking one defect into two and crediting the §11.4.55 signal to the WRONG item), then LINK it, and REOPEN (§11.4.34) **if the resolved head is in a TERMINAL state**; multi-match $\Rightarrow$ resolve every match to its head and act on the SINGLE survivor, >1 distinct head or a broken/circular link $\Rightarrow$ UNDECIDED (ask / mint-with-link, never guess §11.4.6). **Open-original case:** a matched item still OPEN gets LINK-ONLY — ‘reopen an open item’ is undefined in the §11.4.15 status machine and §11.4.34's `Reopened` presupposes a closure to demote FROM (§11.4.7); the reopen path exists to correct the specific lie ‘closed as fixed and it is not’, and an open item is not telling it. **§11.4.202 precedence (explicit, never left to a consumer):** §11.4.202's mandate is that no report is ever LOST — SATISFIED IN FULL by a SAME-DEFECT report landing as link + (terminal-only) reopen; it forbids the report evaporating into prose, it does NOT require a fresh id per report. §11.4.214 refines WHICH item a report lands on;

§11.4.202 owns THAT it lands; DISTINCT/UNDECIDED verdicts run §11.4.202's mint path unchanged. (1) EVERY intake path binds (reporting directives §11.4.202, AI-detected test/gate failures, regression-guard verdicts, cycle re-discovery, operator manual testing — §11.4.34's own reason vocabulary enumerates them; wiring dedup ONLY into the reporting channel under-covers exactly the machine-driven paths that recur most). (2) Key on ticket else normalised (subject, scope) — NEVER a bare subject substring (§11.4.186). (3) **A wrong merge silently DELETES a real defect — so distinct-but-similar is a FIRST-CLASS verdict:** SAME-DEFECT ⇒ link+reopen, DISTINCT ⇒ mint (correct + expected), UNDECIDED ⇒ ask (§11.4.66/.105) or mint-WITH-a-candidate-duplicate-link; autonomous default (§11.4.101) is mint-with-link over merge — a spurious id is cheap + recoverable, a wrongly-merged defect is LOST (the asymmetry decides). (4) Self-validated oracle (§11.4.107(10)/§11.4.186): golden-good + golden-bad + a **negative-control** (two genuinely DISTINCT same-subject items that MUST NOT merge) — an oracle without it is a false-merge engine. (5) Id stability preserved (§11.4.54 — reopen the original, never renumber/reuse/retire; an already-minted duplicate is LINKED via §11.4.90 duplicate-of, never silently deleted). (6) The restored signal IS the point — an undercounting reopen counter is a §11.4/§11.4.1 bluff at the PRIORITISATION layer: it reports the fragile set as stable and aims the deepest scrutiny at the wrong items. Honest boundary §11.4.6 — dedup makes the signal TRUE, not the item correct, and a same-symptom recurrence with a different root cause is real (§11.4.102 decides; the linked+reopened item is where that is recorded, never a reason to fork). Gate CM-RECURRENCE-LINKS-NOT-MINTS (every CONFIRMED-duplicate link whose target was TERMINAL at link time has a reopen event on that target) — **its own scope is §11.4.201-bound:** it MUST NOT fire on an OPEN-target link (LINK-only is correct there) nor on a clause-(3) UNDECIDED candidate-duplicate link (un-reopened BY DESIGN, and the mandated autonomous default), or it becomes exactly the §11.4.201(1) false-positive refusal this same round harvests against; its golden-FALSE set MUST include both + paired §1.1 mutation. Why NEW: §11.4.55 CONSUMES the count, §11.4.34 presupposes the status is already Reopened , §11.4.54 forbids only one-id-two-items, §11.4.186's DEDUP needs DIVERGENT metadata and gates export/sync/commit not INTAKE, and §11.4.202's operative text, read WITHOUT this anchor, reads as mint-every-time — the reading the precedence paragraph above corrects (they are NOT in conflict: §11.4.202 owns THAT a report lands, §11.4.214 refines WHICH item it lands on).

§11.4.215 (NEW) — A doc that BINDS work MUST live, tracked, in the repository where that work happens. FACT: the document DEFINING a binding acceptance gate was UNTRACKED and lived in an out-of-tree scratch directory belonging to a different checkout — so the gate could not be honoured (unreadable), reviewed (no diff/history/authorship), version-controlled (no revision/rollback), or reconciled when it contradicted the source of truth. A binding artifact that lives nowhere is a rumour with an imperative mood. (1) **BINDING ⇒ TRACKED, and un-tracked ⇒ NOT BINDING** — a binding claim sourced from an untracked/out-of-tree doc MUST NOT be enforced; land the doc or decline the gate, never obey a citation of it. (2) **IN THE REPOSITORY THE WORK HAPPENS IN** — a sibling checkout / scratch dir / chat message / issue comment is NOT tracked for this purpose (the test: can someone who clones ONLY that repo read the binding text at its declared path?); multi-repo binding docs are inherited BY REFERENCE from a tracked shared dependency (§11.4.28/§11.4.177), never copied, never out-of-tree. (3) **The prior question §11.4.212 cannot ask** — §11.4.212 ENUMERATES the in-repo scope and checks reachability, so a doc in no checkout is never enumerated and its gate cannot FAIL; §11.4.212 answers “is it findable?”, §11.4.215 answers “is it here at all?” (once in-repo, §11.4.212/.44/.65/.73/.106 all attach as normal — this is their precondition, not their competitor). (4) **The generalisation of §11.4.95** (whose entire content is “authoritative ⇒ tracked, NOT a build artefact, IS authoritative source data” for ONE named artifact), exactly as §11.4.212 was minted as the generalisation of §11.4.57. §11.4.11 is the counterweight and stays intact — logs/forensics/scratch remain out-of-tree and untracked by DEFAULT; **binding is what pulls a document into the repository**, and a doc nobody cites as binding is governed by §11.4.11, not this anchor. Honest boundary §11.4.6 — in-repo presence proves AVAILABLE + ATTRIBUTABLE + VERSION-CONTROLLED, never CORRECT/current (§11.4.44/.106/.186), and never well-seated (§11.4.120’s seam clause — whose forensic case was carried by this SAME out-of-tree document: two defects, one artifact, which is itself the argument that an unreviewable binding doc is where bad gates hide). Gate `CM-BINDING-DOC-IN-REPO` + paired §1.1 mutation.

All Classification: universal (§11.4.17) — no hardware/vendor/project literal; consumers supply their intake wiring, id scheme, normalisation function, and binding-doc paths as DATA per §11.4.35. Propagation gates `CM-COVENANT-114-214-PROPAGATION` + `CM-COVENANT-114-215-PROPAGATION` (literals `11.4.214` / `11.4.215`); the literals `11.4.201` / `11.4.112` /

11.4.120 / 11.4.69 are unchanged and their propagation gates stay GREEN.
All recommended mechanism gates + paired §1.1 mutations = gate-code, a
separate work item.